

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405941
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Haggerty; Timothy et al.

Multi-region workflow management

Abstract

Techniques for performing parallel processing jobs across regionally-distributed servers may include executing a requested job in parallel across two or more regional cloud-based servers, the status of parallel jobs may be locally monitored using respective local instances of a status-tracking database. The multiple instances may replicate modifications of one instance of the status database to other instances of the status database to synchronize the status of regional jobs. Replicating modifications across the respective regional instances of the status database facilitates the parallel executing of the jobs across regions while reducing or avoiding unrecoverable delays if the tasks/job hangs or otherwise fails.

Inventors: Haggerty; Timothy (Henrico, VA), Nannapaneni; Venu Kumar (Glen Allen, VA), Kancharla; Ramachandra (Glen Allen, VA), Nair; Pravin (Addison, TX), Subhedar; Priyank Chittaranjan (Herndon, VA)

Applicant: Capital One Services, LLC (McLean, VA)

Family ID: 93464183

Assignee: Capital One Services, LLC (McLean, VA)

Appl. No.: 18/317689

Filed: May 15, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240386008 A1	Nov. 21, 2024

Publication Classification

Int. Cl.: G06F16/00 (20190101); G06F16/23 (20190101); G06F16/25 (20190101); G06F16/27 (20190101)

U.S. Cl.:

CPC **G06F16/2365** (20190101); **G06F16/254** (20190101); **G06F16/27** (20190101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10936206	12/2020	Dittia	N/A	G06F 3/0619
2004/0054583	12/2003	Nye, III	705/16	G06Q 30/0268
2013/0246376	12/2012	Padmanabhan	707/694	G06F 16/254
2014/0337071	12/2013	Stiffler	705/7.13	G06F 3/04847
2017/0006135	12/2016	Siebel	N/A	H04L 69/40

OTHER PUBLICATIONS

Lu, “A Data Center Monitoring Method and System”, Dec. 7, 2016, 8 pages. (Year: 2016). cited by examiner

Joshi, et al., “Queues with Redundancy: Latency-Cost Analysis”, <https://www.rle.mit.edu/sia/wp-content/uploads/2015/07/2015-joshi-soljanin-wornell-mama.pdf>, 3 pgs. ACM SIGMETRICS Performance Evaluation Review, vol. 43, Issue 216, pp. 54-56, Sep. 2015. cited by applicant

Gardner, et al. “Reducing Latency via Redundant Requests: Exact Analysis” https://kgardner.people.amherst.edu/papers/gardner_sigmetrics15.pdf, 14 pgs., Jun. 15-19, 2015. cited by applicant

Elliot et al. “Combining Partial Redundancy and Checkpointing for HPC”, <https://arc.b.csc.ncsu.edu/~mueller/ftp/pub/mueller/papers/icdcs12.pdf>, 12 pgs, 2012. cited by applicant

“What is Azure Batch?”, <https://learn.microsoft.com/en-us/azure/batch/batch-technical-overview>, 5 pgs; Dec. 13, 2021. cited by applicant

Primary Examiner: Vu; Bai D

Attorney, Agent or Firm: Banner & Witcoff, Ltd.

Background/Summary

FIELD OF USE

(1) Aspects of the disclosure relate generally to processing of information using multi-region cloud-based services.

BACKGROUND

(2) Conventional relational database management systems (RDBMS) are limited to single-region deployments, and would require extensive reconfigurations for multi-regional configurations to provide multi-regional multi-master datastores. In particular, such conventional systems would require adaption to allow non-sequential Globally Unique Identifiers (GUIDs) as well as

reconfigurations of the database referential integrity.

SUMMARY

(3) The following presents a simplified summary of various aspects described herein. This summary is not an extensive overview, and is not intended to identify key or critical elements or to delineate the scope of the claims. The following summary merely presents some concepts in a simplified form as an introductory prelude to the more detailed description provided below. Corresponding apparatus, systems, and computer-readable media are also within the scope of the disclosure.

(4) Aspects described herein generally improve how companies use cloud-based services. In additional aspects, based on the improvements in how companies use cloud-based services, the companies' customers may also experience benefits in reduced computing delays. Possible use-cases may include lengthy computing tasks that typically result in unrecoverable delays if the tasks/jobs hang or otherwise fail.

(5) One or more aspects are directed to improving the handling of cloud-based processing requests by executing (e.g. running) a requested job in parallel across two or more regional cloud-based servers, while locally monitoring the status of the parallel jobs using respective local instances of a status database. The multiple instances may be configured to replicate modifications of one instance of the status database to other instances of the status database to synchronize the status of regional jobs. By replicating modifications across the respective regional instances of the status database, the method and techniques of the disclosure facilitate the parallel execution of the particular job/tasks across two or more regions without requiring reporting to an external data store that may be subject to communication and uptime interruptions.

(6) Other aspects may include monitoring the parallel jobs to determine if one or more of the regional jobs is to be terminated or permitted to complete. The determination may be made based on the synchronized status of regional jobs of the status database. By providing job redundancy across multiple regions by implementing jobs executed in-parallel across regions while identifying the statuses of the respect regional jobs, unrecoverable delays can be advantageously avoided or reduced while also improving the efficiency of the redundant jobs using the synchronized instances of the status database. Thus, significant processing delays may be avoided and reduce delays encountered by the companies' customers.

(7) A system of one or more computers may be configured to perform particular operations or actions by virtue of having software, firmware, hardware, or a combination of them installed on the system that in operation causes or cause the system to perform the actions. One or more computer programs may be configured to perform particular operations or actions by virtue of including instructions that, when executed by data processing apparatus, cause the apparatus to perform the actions. As such, corresponding apparatus, systems, and computer-readable media are also within the scope of the disclosure.

(8) These features, along with many others, are discussed in greater detail below.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present disclosure is described by way of example and not limited in the accompanying figures in which like reference numerals indicate similar elements and in which:

(2) FIG. 1 shows an example computing device in accordance with one or more aspects described herein.

(3) FIG. 2 shows an example computing environment in which one or more aspects described herein may be implemented.

(4) FIG. 3 shows an example computing environment in which one or more aspects described

herein may be implemented.

(5) FIG. 4 shows a flowchart for a parallel processing of regionally-distributed jobs according to one or more aspects of the disclosure.

(6) FIG. 5 shows a flowchart of a method for performing a processing job according to one or more aspects of the disclosure.

DETAILED DESCRIPTION

(7) In the following description of the various embodiments, reference is made to the accompanying drawings, which form a part hereof, and in which is shown by way of illustration various embodiments in which aspects of the disclosure may be practiced. It is to be understood that other embodiments may be utilized and structural and functional modifications may be made without departing from the scope of the present disclosure. Aspects of the disclosure are capable of other embodiments and of being practiced or being carried out in various ways. In addition, it is to be understood that the phraseology and terminology used herein are for the purpose of description and should not be regarded as limiting. Rather, the phrases and terms used herein are to be given their broadest interpretation and meaning.

(8) By way of introduction, aspects discussed herein may relate to methods and techniques for improving the handling of cloud-based processing requests by executing a requested job in parallel across two or more regional cloud-based servers, while locally monitoring the status of the parallel jobs using respective local instances of a status database. The multiple instances may be configured to replicate modifications of one instance of the status database to other instances of the status database to synchronize the status of regional jobs. By replicating modifications across the respective regional instances of the status database, the method and techniques of the disclosure facilitate the parallel execution of the particular job/tasks across two or more regions without requiring reporting to an external data store that may be subject to communication and uptime interruptions.

(9) Other aspects may include monitoring the parallel jobs to determine if one or more of the regional jobs is to be terminated or permitted to complete. The determination may be made based on the synchronize the status of regional jobs of the status database. By providing job redundancy across multiple regions by implementing jobs executing in-parallel across regions while identifying the statuses of the respect regional jobs, unrecoverable delays can be advantageously avoided or reduced while also improving the efficiency of the redundant jobs using the synchronized instances of the status database. Thus, significant processing delays may be avoided and reduce delays encountered by the companies' customers.

(10) Before discussing the concepts of the disclosure in greater detail, however, several examples of a computing device that may be used in implementing and/or otherwise providing various aspects of the disclosure will first be discussed with respect to FIG. 1. FIG. 1 illustrates one example of a computing device **101** that may be used to implement one or more illustrative aspects discussed herein. For example, the computing device **101** may, in some embodiments, implement one or more aspects of the disclosure by reading and/or executing instructions and performing one or more actions based on the instructions. In some embodiments, the computing device **101** may represent, be incorporated in, and/or include various devices such as a desktop computer, a computer server, a mobile device (e.g., a laptop computer, a tablet computer, a smart phone, any other types of mobile computing devices, and the like), and/or any other type of data processing device. In an exemplary embodiment, the computing device **101** may be embodied as the computing device **203**, server **201**, server **202**, and/or server **210** as discussed below with reference to FIG. 2.

(11) The computing device **101** may, in some embodiments, operate in a standalone environment. In others, the computing device **101** may operate in a networked environment. As shown in FIG. 1, various network nodes **101**, **105**, **107**, and **109** may be interconnected via a network **103**, such as the Internet. Other networks may also or alternatively be used, including private intranets,

corporate networks, LANs, wireless networks, personal networks (PAN), and the like. Network **103** is for illustration purposes and may be replaced with fewer or additional computer networks. A local area network (LAN) may have one or more of any known LAN topologies and may use one or more of a variety of different protocols, such as Ethernet. Devices **101**, **105**, **107**, **109**, and other devices (not shown) may be connected to one or more of the networks via twisted pair wires, coaxial cable, fiber optics, radio waves, or other communication media. Additionally, or alternatively, the computing device **101** and/or the network nodes **105**, **107**, and **109** may be a server hosting one or more databases. Databases may include, but are not limited to relational databases, non-relational databases, hierarchical databases, distributed databases, in-memory databases, flat file databases, XML databases, NoSQL databases, graph databases, and/or a combination thereof.

(12) As seen in FIG. 1, the computing device **101** may include a processor **111**, RAM **113**, ROM **115**, network interface **117**, input/output interfaces **119** (e.g., keyboard, mouse, display, printer, etc.), and memory **121**. Processor **111** may include one or more computer processing units (CPUs), graphical processing units (GPUs), and/or other processing units such as a processor adapted to perform computations associated with database operations. Input/output **119** may include a variety of interface units and drives for reading, writing, displaying, and/or printing data or files. Input/output **119** may be coupled with a display such as display **120**. Memory **121** may store software for configuring computing device **101** into a special purpose computing device in order to perform one or more of the various functions discussed herein. Memory **121** may store operating system software **123** for controlling overall operation of the computing device **101**, control logic **125** for instructing the computing device **101** to perform aspects discussed herein, database creation and manipulation software **127** and other applications **129**. Control logic **125** may be incorporated in and may be a part of database creation and manipulation software **127**. In other embodiments, the computing device **101** may include two or more of any and/or all of these components (e.g., two or more processors, two or more memories, etc.) and/or other components and/or subsystems not illustrated here.

(13) Devices **105**, **107**, **109** may have similar or different architecture as described with respect to the computing device **101**. Those of skill in the art will appreciate that the functionality of the computing device **101** (or device **105**, **107**, **109**) as described herein may be spread across multiple data processing devices, for example, to distribute processing load across multiple computers, to segregate transactions based on geographic location, user access level, quality of service (QOS), etc. For example, devices **101**, **105**, **107**, **109**, and others may operate in concert to provide parallel computing features in support of the operation of control logic **125** and/or software **127**.

(14) The data transferred to and from various computing devices may include secure and sensitive data, such as confidential documents, customer personally identifiable information, and account data. Therefore, it may be desirable to protect transmissions of such data using secure network protocols and encryption, and/or to protect the integrity of the data when stored on the various computing devices. For example, a file-based integration scheme or a service-based integration scheme may be utilized for transmitting data between the various computing devices. Data may be transmitted using various network communication protocols. Secure data transmission protocols and/or encryption may be used in file transfers to protect the integrity of the data, for example, File Transfer Protocol (FTP), Secure File Transfer Protocol (SFTP), and/or Pretty Good Privacy (PGP) encryption. In many embodiments, one or more web services may be implemented within the various computing devices. Web services may be accessed by authorized external devices and users to support input, extraction, and manipulation of data between the various computing devices in the system **100**. Web services built to support a personalized display system may be cross-domain and/or cross-platform, and may be built for enterprise use. Data may be transmitted using the Secure Sockets Layer (SSL) or Transport Layer Security (TLS) protocol to provide secure connections between the computing devices. Web services may be implemented using the WS-

Security standard, providing for secure SOAP messages using XML encryption. Specialized hardware may be used to provide secure web services. For example, secure network appliances may include built-in features such as hardware-accelerated SSL and HTTPS, WS-Security, and/or firewalls. Such specialized hardware may be installed and configured in the system **100** in front of one or more computing devices such that any external devices may communicate directly with the specialized hardware.

(15) One or more aspects discussed herein may be embodied in computer-usable or readable data and/or computer-executable instructions, such as in one or more program modules, executed by one or more computers or other devices as described herein. Generally, program modules include routines, programs, objects, components, data structures, etc. that perform particular tasks or implement particular abstract data types when executed by a processor in a computer or other device. The modules may be written in a source code programming language that is subsequently compiled for execution, or may be written in a scripting language such as (but not limited to) Python, JavaScript, or an equivalent thereof. The computer executable instructions may be stored on a computer readable medium such as a hard disk, optical disk, removable storage media, solid-state memory, RAM, etc. As will be appreciated by one of skill in the art, the functionality of the program modules may be combined or distributed as desired in various embodiments. In addition, the functionality may be embodied in whole or in part in firmware or hardware equivalents such as integrated circuits, field programmable gate arrays (FPGA), and the like. Particular data structures may be used to more effectively implement one or more aspects discussed herein, and such data structures are contemplated within the scope of computer executable instructions and computer-usable data described herein. Various aspects discussed herein may be embodied as a method, a computing device, a data processing system, or a computer program product. Having discussed several examples of computing devices which may be used to implement some aspects as discussed further below, discussion will now turn to a method for classifying textual data using a machine-learning classifier.

(16) FIG. 2 is a block diagram of an environment in which systems and/or methods described herein may be implemented. As shown in FIG. 2, the environment may include servers **201**, **202**, and **210**, and a computing device **203** connected by a network **204**. The devices, servers, and network may be interconnected via wired connections, wireless connections, or a combination of wired and wireless connections. The server **201**, server **202**, and server **210** may be located in different locations, such as in different regions. The servers **201** and **202** may be configured to parallelly perform one or more jobs or tasks to provide redundant processing of data as discussed herein. Although three servers and one computing device are illustrated, the disclosure is not limited thereto, and the environment may include more or fewer servers and/or devices. In an example, server **210** may be configured to perform job assignments (e.g., job assignments **304** in FIG. 3) to servers **201** and **202**, and/or function as an external data store (e.g., data store **305** in FIG. 3).

(17) The network **204** may include one or more wired and/or wireless networks. For example, network **204** may include a cellular network (e.g., a long-term evolution (LTE) network, a code division multiple access (CDMA) network, a 3G network, a 4G network, a 5G network, another type of next generation network, etc.), a public land mobile network (PLMN), a local area network (LAN), a wide area network (WAN), a metropolitan area network (MAN), a telephone network (e.g., the Public Switched Telephone Network (PSTN)), a private network, an ad hoc network, an intranet, the Internet, a fiber optic-based network, a cloud computing network, or the like, and/or a combination of these or other types of networks.

(18) The number and arrangement of devices and networks shown in FIG. 2 are provided as an example. In practice, there may be additional devices and/or networks, fewer devices and/or networks, different devices and/or networks, or differently arranged devices and/or networks than those shown in FIG. 2. Furthermore, two or more servers shown in FIG. 2 may be implemented

within a single server, or a single server shown in FIG. 2 may be implemented as multiple, distributed servers or in a cloud-based computing environment. For example, a single server may be implemented as multiple servers distributed among several regions. A set of devices of the environment may perform one or more functions described as being performed by another set of devices of the environment. Network **204** may be represented as a single network but may comprise combinations of other networks or subnetworks.

(19) Storage **208** of server **202** may include a cluster with one or more compute instances. Each compute instance may include one or more tasks. Each task may be performed by a container. Other relationships between compute instances, tasks, and container may be used.

(20) The server **210** may store data in storage **212** to be processed by job applications executed (e.g. performed) by servers **201** and/or **202**. The data may be processed in parallel by processor **205** of server **201** and/or processor **207** of server **202**. Additionally, or alternatively, the computing device **203** may store, in storage **210**, data to be processed by job applications. The data from device **203** may be processed in parallel by processor **205** of server **201** and/or processor **207** of server **202**. In this example, the storage **210** may also function as a data store, and the data store may be distributed to have a local instance (e.g. **312**, **313** in FIG. 3) in one or more of the servers **201**, **202**.

(21) Regionally-Distributed Parallel Processing Jobs

(22) FIG. 3 depicts an environment in which a processing job may be parallelly executed across regionally distributed servers. FIG. 3 includes a first region **301** (Region A) and a second region **302** (Region B). In an exemplary embodiment, first region **301** and second region **302** may be implemented by server **201** and server **202**, respectively. Alternatively, server **201** or **202** may be implemented as multiple, distributed servers or cloud-based computing environments, where the regions **301** and **302** are implemented by respective distributed servers/computed environments.

(23) The regions **301** and **302** may receive data **303** to be processed (e.g. extracted) and one or more job assignments **304**. The data **303** and job assignment(s) **304** may be provided by, for example, computing device **203** or another server. In this example, the data **303** and job assignment(s) **304** may be stored in storage **210** and retrieved by processor **209** so as to provide the data **303** and job assignment(s) **304** to the regions **301** and **302**. The environment may also include a data store **305**, which may be implemented by the storage **210** in this example. In an exemplary embodiment, the data store **305** may be comprised of multiple instances that are distributed as a local data store instance **312**, **313** in the regions **301**, **302**, respectively. In this example, the data store instances **312**, **313** may be synchronized across the regions **301**, **302**.

(24) The first region **301** may include a job processing module **306** configured to execute a processing job on data **303** that is received or otherwise available to the first region **301**. The job processing module **306** may be implemented by, for example, processor **205**. Additionally, or alternatively, the data to be processed may be received from the data store **305** (regional instance **312**). The processing job may be provided in the received job assignment **304**. The processing job may include a process identification, one or more validation rules, and/or machine-readable instructions for executing a process on the data **303**. The data **303** may be included in the processing job in one or more aspects, or be separately provided to the first region **301**. Processed data resulting from the processing job may be stored in the data store **305** (regional instance **312** implemented by storage **206**) and/or one or more other internal and/or external storage devices available to first region **301**.

(25) In an exemplary embodiment, the processing job may be an extract, transform, and load (ETL) process that includes an extraction of data, transformation of data, and/or loading of data, but is not limited thereto. For example, the ETL process can be used to extract data (e.g. from legacy systems), cleanse the data to improve data quality and establish consistency, and load data into a target database. In other aspects, the processing job is an extract, load, transform (ELT) process.

(26) The extraction of data may include the copying or exporting of raw data from source locations

to a staging area. The data sources may be structured or unstructured and may include, without limitation, Structured Query Language (SQL) or NoSQL servers, flat files, Email, customer relationship management (CRM) systems, enterprise resource planning (ERP) systems, web pages, etc.

(27) During the transformation, the raw data undergoes data processing in the staging area. The data may be transformed and/or consolidated based on the use case. The transformation of the data may include (but is not limited to) filtering, de-duplicating, cleansing, validating, and/or authenticating the data. Additionally, or alternatively, the processing may include: performing calculations, translations, and/or summarizations based on the raw data (e.g. modifying row and/or column headers for consistency, converting units of measurement, editing text strings, etc.); data audits for quality and/or compliance checks; encrypting and/or protecting data; and/or formatting or otherwise modifying the data (e.g. into tables or joined tables) to match the schema of the target. The loading process may include moving the transformed data from the staging area into a target.

(28) The first region **301** may also include a local instance **310** of a status-tracking database **308**. The status-tracking database **308** may include one or more database models, such as relational databases, non-relational databases, hierarchical databases, distributed databases, in-memory databases, flat file databases, XML databases, NoSQL databases, graph databases, and/or a combination thereof.

(29) In an exemplary embodiment, the status-tracking database **308** may be implemented as regional (local) instances that may be configured to synchronize with other regional instances. For example, first region **301** may include local instance **310** of a status-tracking database **308** while second region **302** includes local instance **311**. The synchronization of the regional instances **310**, **311** forms status-tracking database **308** as shown in FIG. 3. In this example, the status-tracking database **308** may be a “virtual” database that is a collective representation of each of the synchronized regional instances **310**, **311**.

(30) By the synchronization of the regional instances **310**, **311**, modifications of one instance of the status-tracking database **308** is replicated to other instances of the status-tracking database **308** to provide the status of a regional job to other regions. With the synchronization of the regional instances **310**, **311**, the present disclosure advantageously provides that the job/tasks may be parallelly executed across the regionally distributed servers **301**, **302** without requiring reporting to an external data store (e.g., data store **305**), which may be subject to communication and uptime interruptions. Utilizing the status-tracking database **308**, the present disclosure can leverage the parallel execution of the processing jobs to advantageously avoid unrecoverable delays in situations where computing tasks/jobs hang or otherwise fail.

(31) The status-tracking database **308** is configured to store status tracking data **309**. The status tracking data **309** may include regional job statuses for one or more regional processing jobs that are synchronized from respective status-tracking regional instances **310**, **311**. In the example illustrated in FIG. 3, the status of “Job 0” is “in progress” across all regions. “Job 1” was initially completed by second region **302**, followed by first region **301**. “Job 2” as initially completed by first region **301**, while the processing job was not completed (or the status was not made available) by second region **302**. In this example, by providing job redundancy across multiple regions **301** and **302**, the unrecoverable delay that would have resulted from the failure of the processing job by second region **302** is advantageously avoided.

(32) In an exemplary embodiment, the synchronized statuses of the regional jobs may be used to control the executing of processing jobs. Other aspects may include monitoring the parallel jobs to determine if one or more of the regional jobs is to be terminated or permitted to complete. The determination may be based on the synchronized statuses of regional jobs of the status database **308**. For example, a determination can be made to terminate a processing job based on a status of the job being executed by another region as reflected by that other region's local instance of the status-tracking database **308**. In this example, if one of the redundant processing jobs has

completed, and the completed status is reflected in the status-tracking database **308**, the other redundant processing jobs that are still executing may be terminated (or may be allowed to finish if desired). The processor of the respective server may be configured to control or otherwise instruct the other servers executing the redundant jobs to terminate their respective jobs. In an exemplary embodiment, the job processing module **306** may be configured to update the status tracking instance **310** with a completed status upon completion of the processing job. This completed status is then synchronized across the other regional instances. The corresponding job processing modules (e.g., job processing module **307**) may then decide to terminate based on the completed status by the other regional server.

(33) In an exemplary embodiment, the first region **301** may include Lambda function module **314** configured to perform one or more Lambda functions. The Lambda function module **314** may be implemented by, for example, processor **205**.

(34) Lambda functions may include functions that may be used by higher-order functions, such as those functions performed by job processing module **306**. Lambda functions may be used by other Lambda functions and/or used by non-Lambda functions to perform complex processing operations in first region **301**. For example, the Lambda function module **314** may be configured to monitor the processing of jobs and/or tasks performed by job processing module **306** and to write a corresponding status to the regional instances **310** of the status-tracking database **308** and/or otherwise update one or more statuses within the regional instances **310** of the status-tracking database **308**. Additionally, or alternatively, the Lambda function module **314** may receive a status notification from the job processing module **306** that indicates the status (e.g., in progress, job completed, job failed, etc.) of the current processing job(s) being executed by job processing module **306**, and may write the status to the regional instances **310** of the status-tracking database **308** and/or otherwise update one or more statuses within the regional instances **310** of the status-tracking database **308**.

(35) As explained above, the Lambda function module **314** may monitor the status of the processing job being executed by the job processing module **306**. When the Lambda function module **314** determines that the processing job has completed, the Lambda function module **314** may update the regional instance **310** of the status-tracking database **308**. The Lambda function module **315** of the second region **302** may monitor their respective regional instance **311** and detect the completed status by first region **301**. The Lambda function module **315** may then notify the job processing module **307** of the completed job. The job processing module **307** may then decide to terminate the processing job, for example, based on job processing module **306** completing processing of the processing job. In one aspect, the Lambda function module **315** may be configured to control the job processing module **307** to terminate the processing job based on the monitored statuses of the regional instance **311**.

(36) In an exemplary embodiment, the Lambda function module **314** may additionally or alternatively be configured to instruct or request the regional instance **310** of the status-tracking database **308** to synchronize one or more statuses stored therein with one or more other regional instances (e.g., regional instance **311**) of the status-tracking database **308**. The Lambda function module **314** may additionally or alternatively be configured to perform one or more functions for processing the data **303** under the direction of job processing module **306** and/or to otherwise assist in the processing of the data **303** by job processing module **306**.

(37) Like first region **301**, second region **302** may include a job processing module **307** and a regional instance **311** of the status-tracking database **308**. Job processing module **307** may be configured to perform functions similar to job processing module **306**, and discussion of the similar functions has been omitted for brevity.

(38) Second region **302** may additionally include Lambda function module **315**. Lambda function module **315** may be configured to perform one or more Lambda functions similar to the operation of Lambda function **312**. Discussion of these similar operations has been omitted for brevity. As

will be appreciated, Second region **302** may include additional components or omit one or more components found in First region **301**.

(39) Parallel Processing Operations for Regionally-Distributed Jobs

(40) The parallel processing of data using regionally-distributed jobs will be described in further detail with reference to FIGS. **4** and **5**.

(41) FIG. **4** is a flowchart for parallel processing of regionally-distributed jobs that may include the synchronization of regional instances of a status-tracking database **308**. Some or all of the steps of process **400** may be performed using one or more computing devices as described herein, such as servers **201** and **202**.

(42) The parallel operations of process **400** begins at operations **401** and **402**, and transitions to operations **403** and **404**, respectively, where the processing job is received. The received processing job may include the data (e.g., data **303**) to be processed, or the data (e.g., data **303**) may be received separately. The received processing job may be provided in a job assignment, such as job assignment **304**. The processing job may include a process identification, one or more validation rules, and/or machine-readable instructions for performing a process on the data. Operations **403** and **404** may be performed sequentially, at least partially concurrently, concurrently, or simultaneously.

(43) At operations **405** and **406**, preliminary operations are performed. The preliminary operations may include checking if dependencies have successfully completed and/or if dependent libraries are available. If the dependency check fails, the processing job can be retried and/or a notification of the failure can be generated. The preliminary operations may additionally or alternatively include setting up of the environment for the job. This may include building or reusing infrastructure. In an exemplary embodiment, the preliminary operations **405** and/or **406** may be performed before the processing job is received. Operations **405** and **406** may be performed sequentially, at least partially concurrently, concurrently, or simultaneously.

(44) At operations **407** and **408**, the processing job is executed (e.g., run). The parallelly execution of the jobs in different regions can provide redundant processing of data. In one or more aspects, the processing jobs may be an extract, transform, and load (ETL) process, but is not limited thereto.

(45) Operations **407** and **408** across the regions are performed concurrently. The operations **407** and **408** may be performed at least partially concurrently. However, the greater amount of time that the operations **407** and **408** are being performed concurrently, the more process **400** may advantageously reduce and/or avoid potential delays if the jobs hang or otherwise fail. In an exemplary embodiment, the operations **407** and **408** are performed simultaneously.

(46) During execution of the processing jobs across the regions, the status of the executing jobs may be monitored and used to update a status-tracking database **308** (operation **413**). The status-tracking database **308** may be implemented as regional (local) instances within the respective regions. The regional instance (e.g. **310**) may be configured to synchronize with other regional instances (e.g. **311** of second region **302**). The synchronization of the regional instances **310**, **311** forms status-tracking database **308** as shown in FIG. **3**. By the synchronization of the regional instances **310**, **311**, modifications of one instance of the status-tracking database **308** is replicated to other instances of the status-tracking database **308** to provide the status of a regional job to other regions. With the synchronization of the regional instances **310**, **311**, the present disclosure advantageously provides that the job/tasks may be executed in-parallel across the regionally distributed servers **301**, **302** without requiring reporting to an external data store (e.g. data store **305**) that may be subject to communication and uptime interruptions. Further, the parallel execution advantageously avoids or reduces unrecoverable delays in situations where computing tasks/jobs hang or otherwise fail.

(47) In an exemplary embodiment, for the status tracking at operation **413**, the synchronized statuses of the regional jobs may be used to control the executing processing jobs. Other aspects may include monitoring the parallel jobs to determine if one or more of the regional jobs is to be

terminated or permitted to complete. The determination may be based on the synchronized statuses of regional jobs for the status database **308**. For example, a determination can be made to terminate a processing job based on a status of another processing job. In this example, if one of the redundant processing jobs has completed, and the completed status is reflected in the status-tracking database, the other redundant processing jobs that are still executing may be terminated. The processor of the respective server may be configured to control or otherwise instruct the other servers executing the redundant jobs to terminate their respective jobs.

(48) In an exemplary embodiment, the job processing module **306** may be configured to update the status tracking instance **310** with a completed status upon completion of the processing job. This completed status is then synchronized across the other regional instances. The corresponding job processing modules may then decide to terminate based on the completed status by the other regional server. In one or more aspects, the Lambda function module **314** may monitor the status of the processing job being executed by the job processing module **306**. When the Lambda function module **314** determines that the processing job has completed, the Lambda function module **314** may update the regional instance **310** of the status-tracking database **308**. The Lambda function module **315** of the other region(s) may monitor their respective regional instance **311** of the second region **302** and detect the completed status by first region **301**. The Lambda function module **315** may then notify the job processing module **307** that job processing module **306** has completed the job. The job processing module **307** may then decide to terminate the processing job. In one aspect, the Lambda function module **315** may be configured to control the job processing module **307** to terminate the processing job based on the monitored statuses of the regional instance **311**.

(49) At operations **409** and/or **410**, the processed data may be validated. The processed data is generated by the completed processing job, and operation **409** and/or **410** may be skipped/omitted if the processing jobs executing in operations **407** and/or **408** are terminated, stalled, or otherwise fail. Operations **409** and **410** may be performed sequentially, at least partially concurrently, concurrently, or simultaneously. Operation **409** and/or operation **410** may be omitted or skipped if the respective jobs at operations **407** and **408** is halted or otherwise fails.

(50) At operations **411** and **412**, the process **400** ends. The process **400** in one or more of the regions may end before operations **411**, **412**, for example, based on the status of the processing jobs performed by operations **407** and **408**.

(51) The processing of data using a processing job that may be redundantly executed across one or more regions will be described in further detail with reference to FIG. 5. Discussion of process **500** of FIG. 5 may reference FIG. 3 and the process **400** of FIG. 4. FIG. 5 is a flowchart for a executing (performing) a processing job. The process **500** includes the operations from the perspective of a server executing a received processing job. Some or all of the steps of process **500** may be performed using one or more computing devices as described herein, such as server **201** and/or server **202**.

(52) The process **500** begins as operation **501**, and transitions to operation **503**, where the processing job and data to be processed by the processing job is received (e.g. by servers **201** and **202** from server **210**). The received processing job may include the data to be processed, or the data may be received separately. The received processing job may be provided in a job assignment, such as job assignment **304**. The processing job may include a process identification, one or more validation rules, and/or machine-readable instructions for performing a process on the data **303**.

(53) The process **500** transitions to operation **505**, where the processing job is executed. In this example, a sever (e.g. first server **201**, second server, n.sup.th server **210**) may execute the processing job on the received data to generate processed data. The execution of the processing job may include reading the received data, processing the data to generate processed data, and/or validating the processed data (e.g. based on validation rules included in the job assignment or otherwise separately provided to the server). The processing jobs may be an extract, transform, and load (ETL) process, but is not limited thereto. For example, the processing job may be an

extraction sub-process of an ETL process; an extraction sub-process and a transform sub-process of an ETL process; or an extraction sub-process and a load sub-process of an extract, load, and transform (ELT) process.

(54) After operation **505**, the process **500** transitions to operation **507**, where a status of the processing job is checked to determine if the processing job has completed. If the job has completed (YES at operation **507**), which may include a status of “completed” or the like, the process **500** transitions to operation **509**. If the job has not completed (NO at operation **507**), which may include a status of “processing,” “in progress,” “active,” or the like, the process **500** transitions to operation **511**.

(55) At operation **509**, the local instance **310** of the status-tracking database **308** is updated to indicate that the processing job has completed. The “completed” status is then populated to the other local instances (e.g. **311**) of the status-tracking database **308** to synchronize the job status across the various regions. The one or more other regions parallelly executing the job may terminate the active job or otherwise allow it to complete. After operation **509**, the process **500** transitions to operation **519** where the process **500** ends. The process **500** may be repeated for the next processing job request.

(56) At operation **511**, the local instance **310** of the status-tracking database **308** is updated to indicate that the processing job has not completed. This “in progress” status is then populated to the other local instances (e.g. **311**) of the status-tracking database **308** to synchronize the job status across the various regions. The synchronization of the “in progress” status can be used by the other regions and inform the other regions to continue their active processing job. If the check returns a status indicating that the job is has not completed and is in progress, it may be determined that the processing job is stalled (e.g., hung), for example, if a predetermined time threshold/limit has been exceeded.

(57) After operation **511**, the process transitions to operation **513**, where the local instance **310** of the tracking status database **308** is queried to determine the status of the processing jobs in the other region(s). The query can include a request to the local instance **310** (e.g. by the Lambda function module **314**) for a status report, which can then be provided by the local instance **310**. In other aspects, the status from the local instance **310** can be read/pulled (e.g. by the Lambda function module **314**) without submitting a request for the information.

(58) After operation **513**, the process **500** transitions to operation **515**, where it is determined if the processing job has been completed by one or more other regional servers. The determination can be based on the status obtained from the local instance **310** in operation **513**. If the obtained status indicates that the processing job has not been completed (NO at operation **515**) by one or more other regional servers, the process **500** returns to operation **507**. If the obtained status indicates that the processing job has been completed (YES operation **515**) by one or more other regional servers, the process **500** transitions to operation **517**.

(59) At operation **517**, the execution of the processing job by a sever (e.g. first server **201**, second server, n.sup.th server **210**) may be terminated. In one or more aspects, the processing job can be allowed to complete instead of being terminated. After operation **517**, the process **500** transitions to operation **519** where the process **500** ends. The process **500** may be repeated for the next processing job request.

(60) One or more aspects discussed herein may be embodied in computer-usable or readable data and/or computer-executable instructions, such as in one or more program modules, executed by one or more computers or other devices as described herein. Generally, program modules include routines, programs, objects, components, data structures, and the like. that perform particular tasks or implement particular abstract data types when executed by a processor in a computer or other device. The modules may be written in a source code programming language that is subsequently compiled for execution, or may be written in a scripting language such as (but not limited to) Python, Perl, or any other suitable scripting language. The computer executable instructions may be

stored on a computer readable medium such as a hard disk, optical disk, removable storage media, solid-state memory, RAM, and the like. As will be appreciated by one of skill in the art, the functionality of the program modules may be combined or distributed as desired in various embodiments. In addition, the functionality may be embodied in whole or in part in firmware or hardware equivalents such as integrated circuits, field programmable gate arrays (FPGA), and the like. Particular data structures may be used to more effectively implement one or more aspects discussed herein, and such data structures are contemplated within the scope of computer executable instructions and computer-usable data described herein. Various aspects discussed herein may be embodied as a method, a computing device, a system, and/or a computer program product.

(61) Although the present invention has been described in certain specific aspects, many additional modifications and variations would be apparent to those skilled in the art. In particular, any of the various processes described above may be performed in alternative sequences and/or in parallel (on different computing devices) in order to achieve similar results in a manner that is more appropriate to the requirements of a specific application. It is therefore to be understood that the present invention may be practiced otherwise than specifically described without departing from the scope and spirit of the present invention. Thus, embodiments of the present invention should be considered in all respects as illustrative and not restrictive. Accordingly, the scope of the invention should be determined not by the embodiments illustrated, but by the appended claims and their equivalents.

(62) It is to be understood that other embodiments may be utilized and structural and functional modifications may be made without departing from the scope of the present disclosure. Aspects of the disclosure are capable of other embodiments and of being practiced or being carried out in various ways. Also, it is to be understood that the phraseology and terminology used herein are for the purpose of description and should not be regarded as limiting. Rather, the phrases and terms used herein are to be given their broadest interpretation and meaning. The use of “including” and “comprising” and variations thereof is meant to encompass the items listed thereafter and equivalents thereof as well as additional items and equivalents thereof. Any sequence of computer-implementable instructions described in this disclosure may be considered to be an “algorithm” as those instructions are intended to solve one or more classes of problems or to perform one or more computations. While various directional arrows are shown in the figures of this disclosure, the directional arrows are not intended to be limiting to the extent that bi-directional communications are excluded. Rather, the directional arrows are to show a general flow of steps and not the unidirectional movement of information. In the entire specification, when an element is referred to as “comprising” or “including” another element, the element should not be understood as excluding other elements so long as there is no special conflicting description, and the element may include at least one other element. In addition, the terms “unit” and “module”, for example, may refer to a component that exerts at least one function or operation, and may be realized in hardware or software, or may be realized by combination of hardware and software. In addition, terms such as “. . . unit,” “. . . module” described in the specification mean a unit for performing at least one function or operation, which may be implemented as hardware or software, or as a combination of hardware and software. Throughout the specification, expression “at least one of a, b, and c” may include ‘a only,’ ‘b only,’ ‘c only,’ ‘a and b,’ ‘a and c,’ ‘b and c,’ and/or ‘all of a, b, and c.’

(63) It is noted that various connections between elements are discussed in the following description. It is noted that these connections are general and, unless specified otherwise, may be direct or indirect, and that the specification is not intended to be limiting in this respect. As described herein, thresholds are referred to as being “satisfied” to generally encompass situations involving thresholds above increasing values as well as encompass situations involving thresholds below decreasing values. The term “satisfied” is used with thresholds to address when values have passed a threshold and then approaching the threshold from an opposite side as using terms such as

“greater than,” “greater than or equal to,” “less than,” and “less than or equal to” can add ambiguity where a value repeated crosses a threshold.

Claims

1. A method comprising: receiving, by a computing device and from a first remote server, a first processing job comprising first data and one or more validation rules; performing, by the computing device, the first processing job resulting in a processed form of the first data, wherein the performing comprises: reading the first data; validating, based on the one or more validation rules, the first data; determining, by the computing device and via a local instance of a status database of the computing device, that the second remote server has not completed the first processing job; based on the determining that the second remote server has not completed the first processing job: writing, by the computing device, to the local instance of the status database, a status of the first processing job, wherein the status indicates that the computing device as completed the first processing job; transmitting, by the computing device and to the second remote server, an instruction to terminate the first processing job by writing, to the local instance of a status database of the second remote server, the status; and writing, by the computing device, the processed form of the first data based on the performing the first processing job.
2. The method of claim 1, further comprising: receiving, by the computing device from the first remote server, a second processing job comprising second data and one or more second validation rules; performing, by the computing device, the second processing job, wherein the performing comprises: reading the second data; validating, based on the one or more second validation rules, at least some of the second data; determining a failure of the performing of the second processing job by the computing device; and writing, by the computing device to the local instance of the status database of the computing device and based on the failure of the performing of the second processing job by the computing device, a second status of the second processing job, wherein the second status indicates the failure of performing the second processing job by the computing device.
3. The method of claim 1, further comprising: receiving, by the computing device from the first remote server, a second processing job comprising a second process identification, second data, one or more second validation rules, and second machine-readable instructions to perform a second process on the second data; performing, by the computing device, the second processing job, wherein the performing comprises: reading the second data; validating, based on the one or more second validation rules, at least some of the second data; determining, by the computing device and from the local instance of the status database of the computing device, that the second remote server has completed the second processing job; and terminating, by the computing device and based on the determining that the second remote server has completed the second processing job, the performing of the second processing job.
4. The method of claim 1, wherein the first processing job is an extraction sub process of an extraction, transform, and load process.
5. The method of claim 1, wherein determining the status of the first processing job performed by the second remote server comprises: synchronizing the local instance of the status database of the computing device and a local instance of the status database of the second remote server to determine the status of the first processing job performed by the second remote server.
6. The method of claim 1, further comprising executing, by the computing device, a lambda function configured to monitor the status of the first processing job being performed by the computing device and to write the status of the first processing job being performed by the computing device to the local instance of the status database of the computing device.
7. The method of claim 1, wherein the first data comprises data from a single data storage.
8. The method of claim 1, wherein: the first data comprises a combination of second data from a

first data storage and third data from a second data storage, and a format of the second data corresponds to a format of the third data.

9. The method of claim 1, wherein: the first data comprises a combination of second data from a first data storage and third data from a second data storage, a format of the second data differs from a format of the third data.

10. An apparatus comprising: one or more processors; and memory storing instructions that, when executed by the one or more processors, cause the apparatus to: receive, by a computing device and from a first remote server, a first processing job comprising first data and one or more validation rules; perform, by the computing device, the first processing job resulting in a processed form of the first data, wherein the instructions to perform further cause the apparatus to: execute one or more lambda functions; read the first data; and validate, based on the one or more validation rules, the first data; determine, via a local instance of a status database of the computing device, that the second remote server has not completed the first processing job; based on the determination that the second remote server has not completed the first processing job: write, to the local instance of the status database, a status of the first processing job, wherein the status indicates that the computing device has completed the first processing job; transmit, to the second remote server, an instruction to terminate the first processing job by writing, to the local instance of a status database of the second remote server, the status; and write, by the computing device, the processed form of the first data based on the performance of the first processing job.

11. The apparatus of claim 10, wherein the instructions, when executed by the one or more processors, further cause the apparatus to: receive, by the computing device from the first remote server, a second processing job comprising second data and one or more second validation rules; perform, via the computing device, the second processing job, wherein the instructions to perform further cause the apparatus to: read the second data; and validate, based on the one or more second validation rules, at least some of the second data; determine a failure of the performing of the second processing job by the computing device; and write, by the computing device to the local instance of the status database of the computing device and based on the failure of the performing of the second processing job by the computing device, a second status of the second processing job, wherein the second status indicates the failure of performing the second processing job by the computing device.

12. The apparatus of claim 10, wherein the instructions, when executed by the one or more processors, further cause the apparatus to: receive, by the computing device and from the first remote server, a second processing job comprising a second process identification, second data, one or more second validation rules, and second machine-readable instructions to perform a second process on the second data; and perform, via the local computing device, the second processing job, wherein the instructions to perform further cause the apparatus to: read the second data; validate, based on the one or more second validation rules, at least some of the second data; determine, by the computing device and via the local instance of the status database, whether the second remote server has completed the second processing job; and terminate, by the computing device and based on the determining that the second remote server completed the second processing job, the performing of the second processing job.

13. The apparatus of claim 10, wherein the first processing job is an extraction sub process of an extraction, transform, and load process.

14. The apparatus of claim 13, wherein the first processing job further comprises a transform sub process of the extraction, transform, and load process.

15. The apparatus of claim 10, wherein the instructions, when executed by the one or more processors, further cause the apparatus to determine the status of the first processing job performed by the second remote server by synchronizing the local instance of the status database of the computing device and a local instance of the status database of the second remote server to determine the status of the first processing job performed by the second remote server.

16. The apparatus of claim 10, wherein the first data comprises data from a single data storage.
17. The apparatus of claim 10, wherein: the first data comprises a combination of second data from a first data storage and third data from a second data storage, and a format of the second data corresponds to a format of the third data.
18. The apparatus of claim 10, wherein: the first data comprises a combination of second data from a first data storage and third data from a second data storage, and a format of the second data differs from a format of the third data.
19. One or more non-transitory, computer-readable media storing instructions that, when executed by one or more processors, cause the one or more processors to perform steps comprising: receiving, by a computing device and from a first remote server, a first processing job comprising a process identification, first data, one or more validation rules, and machine-readable instructions to perform a process on the first data; performing, by the computing device, the first processing job resulting in a processed form of the first data, wherein the performing comprises: reading the first data; validating, based on the one or more validation rules, the first data; determining, by the computing device and via a local instance of a status database of the computing device, that the second remote server has not completed the first processing job, wherein the local instance of the status database of the computing device is configured to replicate changes to the local instance of the status database of the computing device to a local instance of a status database of the second remote server; based on the determining that the second remote server has not completed the first processing job: writing, by the computing device and to the local instance of the status database, a status of the first processing job, wherein the status indicates that the computing device has completed the first processing job; transmitting, by the computing device and to the second remote server, an instruction to terminate the first processing job by writing, to the local instance of the status database of the second remote server, the status; and writing, by the computing device, the processed form of the first data based on the performing the first processing job.
20. The one or more non-transitory, computer-readable media of claim 19, wherein the instructions cause the one or more processors to perform further steps comprising: receiving, by the computing device and from the first remote server, a second processing job comprising a second process identification, second data, one or more second validation rules, and second machine-readable instructions to perform a second process on the second data; performing, via the computing device, the second processing job, wherein the performing comprises: reading the second data; and validating, based on the one or more second validation rules, at least some of the second data; determining a failure of the performing of the second processing job by the computing device; and writing, by the computing device and to the local instance of the status database of the computing device, based on the failure of the performing of the second processing job by the computing device, a second status of the second processing job, wherein the second status indicates the failure of performing the second processing job by the computing device.
-