

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405949
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Clemens; David et al.

Generating, accessing, and displaying lineage metadata

Abstract

Among other things, we describe a method of receiving a portion of metadata from a data source, the portion of metadata describing nodes and edges; generating instances of a data structure representing the portion of metadata, at least one instance of the data structure including an identification value that identifies a corresponding node, one or more property values representing respective properties of the corresponding node, and one or more pointers to respective identification values, each pointer representing an edge associated with a node identified by the corresponding respective identification value; storing the instances of the data structure in random access memory; receiving a query that includes an identification of at least one particular element of data; and using at least one instance of the data structure to cause a display of a computer system to display a representation of lineage of the particular element of data.

Inventors: Clemens; David (Cambridge, MA), Radivojevic; Dusan (North Andover, MA), Galarneau; Neil (West Roxbury, MA)

Applicant: Ab Initio Technology LLC (Lexington, MA)

Family ID: 61656328

Assignee: Ab Initio Technology LLC (Lexington, MA)

Appl. No.: 18/345706

Filed: June 30, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240078229 A1	Mar. 07, 2024

Related U.S. Application Data

continuation parent-doc US 15829152 20171201 US 11741091 child-doc US 18345706

Publication Classification

Int. Cl.: **G06F16/245** (20190101); **G06F16/22** (20190101); **G06F16/248** (20190101); **G06F16/83** (20190101); **G06F40/117** (20200101)

U.S. Cl.:

CPC **G06F16/245** (20190101); **G06F16/22** (20190101); **G06F16/248** (20190101); **G06F16/83** (20190101); **G06F40/117** (20200101);

Field of Classification Search

CPC: G06F (16/22); G06F (16/245); G06F (16/248); G06F (16/83); G06F (16/81); G06F (16/832); G06F (16/8365)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6003040	12/1998	Mital et al.	N/A	N/A
6725227	12/2003	Li	N/A	N/A
7401064	12/2007	Arone	N/A	N/A
7456840	12/2007	Molesky et al.	N/A	N/A
7493570	12/2008	Bobbin	N/A	N/A
7546226	12/2008	Yeh et al.	N/A	N/A
7590672	12/2008	Slik et al.	N/A	N/A
7725433	12/2009	Labrie	N/A	N/A
7844582	12/2009	Arbilla et al.	N/A	N/A
7970240	12/2010	Chao et al.	N/A	N/A
8266122	12/2011	Newcombe et al.	N/A	N/A
8332782	12/2011	Chang et al.	N/A	N/A
8515911	12/2012	Zhou	707/649	G06F 16/1873
8577852	12/2012	Haahenson et al.	N/A	N/A
8654125	12/2013	Gibson	N/A	N/A
8819010	12/2013	Fankhauser	N/A	N/A
11741091	12/2022	Clemens et al.	N/A	N/A
2002/0030703	12/2001	Robertson et al.	N/A	N/A
2004/0181554	12/2003	Heckerman et al.	N/A	N/A
2004/0255239	12/2003	Bhatt	N/A	N/A
2005/0114369	12/2004	Gould et al.	N/A	N/A
2005/0246352	12/2004	Moore et al.	N/A	N/A
2006/0095466	12/2005	Stevens et al.	N/A	N/A
2006/0106847	12/2005	Eckardt et al.	N/A	N/A
2006/0149782	12/2005	Yeh et al.	N/A	N/A
2006/0190844	12/2005	Binder et al.	N/A	N/A
2006/0218159	12/2005	Murphy et al.	N/A	N/A

2006/0271505	12/2005	Vierich et al.	N/A	N/A
2006/0294150	12/2005	Stanfill et al.	N/A	N/A
2007/0016624	12/2006	Powers	N/A	N/A
2007/0033220	12/2006	Drucker et al.	N/A	N/A
2007/0061287	12/2006	Le et al.	N/A	N/A
2007/0061353	12/2006	Bobbin	N/A	N/A
2007/0112875	12/2006	Yagawa	N/A	N/A
2007/0150496	12/2006	Feinsmith	N/A	N/A
2007/0255741	12/2006	Geiger et al.	N/A	N/A
2008/0040388	12/2007	Petri	N/A	G06F 16/907
2008/0155119	12/2007	Imamura et al.	N/A	N/A
2008/0163124	12/2007	Bonev et al.	N/A	N/A
2008/0172629	12/2007	Tien et al.	N/A	N/A
2008/0183658	12/2007	Mangipudi	N/A	N/A
2008/0244721	12/2007	Barrus	726/9	G06F 21/6254
2009/0012983	12/2008	Senneville et al.	N/A	N/A
2009/0216728	12/2008	Brainerd et al.	N/A	N/A
2009/0224941	12/2008	Kansal et al.	N/A	N/A
2010/0138420	12/2009	Bator et al.	N/A	N/A
2010/0138431	12/2009	Bator et al.	N/A	N/A
2010/0223430	12/2009	Walker et al.	N/A	N/A
2010/0313157	12/2009	Carlsson	N/A	N/A
2011/0320460	12/2010	Fankhauser	N/A	N/A
2012/0059857	12/2011	Jackson, Jr.	N/A	N/A
2012/0143825	12/2011	Boehm et al.	N/A	N/A
2012/0209656	12/2011	Rojahn	N/A	N/A
2012/0254805	12/2011	Pic et al.	N/A	N/A
2012/0310875	12/2011	Prahlad et al.	N/A	N/A
2012/0311496	12/2011	Cao	N/A	N/A
2013/0332423	12/2012	Puri	N/A	N/A
2014/0019423	12/2013	Liensberger	N/A	N/A
2014/0114907	12/2013	Kozina	N/A	N/A
2014/0244687	12/2013	Shmueli et al.	N/A	N/A
2014/0279979	12/2013	Yost et al.	N/A	N/A
2015/0012478	12/2014	Mohammad et al.	N/A	N/A
2015/0142724	12/2014	Nair	N/A	N/A
2015/0271267	12/2014	Solis	709/213	G06F 16/178
2016/0019057	12/2015	Bach et al.	N/A	N/A
2016/0063106	12/2015	Chai	N/A	N/A
2016/0232229	12/2015	Radivojevic	N/A	N/A
2016/0232230	12/2015	Radivojevic	N/A	N/A
2017/0058320	12/2016	Locke et al.	N/A	N/A
2017/0154087	12/2016	McClure	N/A	G06Q 10/0633
2017/0286526	12/2016	Bar-Or et al.	N/A	N/A
2018/0157702	12/2017	Clemens et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
------------	------------------	---------	-----

101084496	12/2006	CN	N/A
101971165	12/2010	CN	N/A
102239458	12/2010	CN	N/A
102890720	12/2012	CN	N/A
1258814	12/2001	EP	N/A
1510937	12/2004	EP	N/A
H08-033895	12/1995	JP	N/A
H11-307412	12/1998	JP	N/A
2002288403	12/2001	JP	N/A
2003044538	12/2002	JP	N/A
2005122703	12/2004	JP	N/A
2006-190261	12/2005	JP	N/A
2008059563	12/2007	JP	N/A
2008134705	12/2007	JP	N/A
2008524671	12/2007	JP	N/A
2010244157	12/2009	JP	N/A
2011517352	12/2010	JP	N/A
2013-077063	12/2012	JP	N/A
2015-531129	12/2014	JP	N/A
2016520890	12/2015	JP	N/A
WO 2001082068	12/2000	WO	N/A
WO 2001082072	12/2000	WO	N/A
WO 2005086906	12/2004	WO	N/A
WO 2007002647	12/2006	WO	N/A
WO 2010065623	12/2009	WO	N/A
WO 2015183738	12/2014	WO	N/A
WO 2016014615	12/2015	WO	N/A

OTHER PUBLICATIONS

Alvarez et al., "A compact representation of graph databases," InProceedings of the Eighth Workshop on Mining and Learning with Graphs, Jul. 24, 2010, 18-25. cited by applicant

Alvarez-Garcia et al., "Compact and efficient representation of general graph databases," Knowledge and Information Systems., Sep. 1, 60(3):1479-510. cited by applicant

Australian Office Action in 2009-322441, issued Dec. 16, 2015, 4 pages. cited by applicant

Chinese Office Action, Application No. 200980148414.6, mailed May 10, 2013, 15 pages (with English translation). cited by applicant

DePauw, W., et al., "Web Services Navigator: Visualizing the Execution of Web Services." IBM Systems Journal, vol. 44, No. 4, 2005, pp. 821-845. cited by applicant

Hiroto Namihara, "Learning Data Structures and Algorithm in C language," Interface, Japan, CQ Publishing Co., Ltd., Apr. 1, 1990, vol. 16, No. 4, pp. 158-181. cited by applicant

International Preliminary Report on Patentability in PCT Appln. No. PCT/US2017/064227, dated Jun. 4, 2019, 14 pages. cited by applicant

International Preliminary Report on Patentability, PCT/US2013/062369, Apr. 9, 2015, 7 pages. cited by applicant

International Search Report & Written Opinion issued in PCT application No. PCT/US09/66390, dated Jan. 20, 2010, 6 pages. cited by applicant

International Search Report & Written Opinion issued in PCT application No. PCT/US09/66394, dated Jan. 20, 2010, 6 pages. cited by applicant

International Search Report & Written Opinion, PCT/US2013/062369, Feb. 28, 2014, 9 pages. cited by applicant

International Search Report & Written Opinion, PCT/US2013/076407, Jun. 18, 2014, 8 pages. cited by applicant

International Search Report and Written Opinion in PCT/US14/26133, Aug. 21, 2014, 8 pages. cited by applicant

Japanese Office Action (with English translation), Application No. 2011-514796, mailed Jun. 20, 2013, 9 pages. cited by applicant

Japanese Office Action (with English translation), Application No. 2011-539651, mailed Oct. 11, 2013, 4 pages. cited by applicant

Japanese Office Action (with English translation), Application No. 2011-539652, mailed Dec. 16, 2013, 8 pages. cited by applicant

Japanese Office Action (with English translation), Application No. 2011-539652, mailed Nov. 14, 2014, 7 pages. cited by applicant

Japanese Office Action (with English translation), Application No. 2014-232962, mailed Nov. 13, 2015, 8 pages. cited by applicant

JP Office Action in Japanese Appln. No. 2019-525760, dated Aug. 28, 2020, 7 pages (with English translation). cited by applicant

Karr, et al., "Data quality: A Statistical Perspective," Statistical Methodology 3 2006, pp. 137-173. cited by applicant

Korean Office Action in Application No. 10-2011-7013820, mailed Oct. 21, 2015, 14 pages. cited by applicant

Korean Office Action in Application No. 10-2011-7014656, mailed Oct. 21, 2015, 14 pages. cited by applicant

Nelson, Theodor Holm, "A Cosmology for a Different Computer Universe: Data Model, Mechanisms, Virtual Machine and Visualization Infrastructure," Journal of Digital Information, vol. 5, No. 1, 22 pages (2004). cited by applicant

Office Action in Chinese Appln. No. 201780074708.3, dated Jan. 5, 2023, 15 pages (with English translation). cited by applicant

Office Action in Chinese Appln. No. 201780074708.3, mailed on Nov. 15, 2023, 23 pages (with Machine translation). cited by applicant

Office Action in European Appln. No. 17851913.8, dated Jun. 14, 2022, 29 pages. cited by applicant

Office Action in Japanese Appln. No. 2019-525760, dated Jul. 26, 2021, 9 pages (with English translation). cited by applicant

Office Action in Japanese Appln. No. 2019-525760, dated Jun. 1, 2022, 22 pages (with English translation). cited by applicant

Office Action in Japanese Appln. No. 2021-191498, dated Feb. 8, 2023, 15 pages (with English translation). cited by applicant

Office Action in Singapore Appln. No. 11201903995V, dated Jun. 23, 2022, 8 pages. cited by applicant

Parker et al., Visualization of Large Nested Graphs in 3D: Navigation and Interaction, University of New Brunswick, Jul. 18, 2000 <http://com.unh.edu/vislab/PDFs/visualnav.pdf>. cited by applicant

Robertson et al., "Polyarchy Visualization: Visualizing Multiple Intersecting Hierarchies," Computer Human Interaction—CHI , pp. 423-430, 2002. cited by applicant

Robertson, George G. "From Hierarchies to Polyarchies: Visualizing Multiple Relationships," Advanced Visual Interfaces Jan. 2000. cited by applicant

SG Office Action in Singapore Appln. No. 11201903995V, dated Oct. 13, 2020, 7 pages. cited by applicant

Supplementary European Search Report in EP14768183 mailed Oct. 4, 2016 (5 pages). cited by applicant

Supplementary European Search Report, EP 09767701, Sep. 25, 2012, 7 pages. cited by applicant

Supplementary European Search Report, EP 09831039, May 15, 2012, 9 pages. cited by applicant
Supplementary European Search Report, EP 09831042, May 4, 2012, 8 pages. cited by applicant
Takahashi, Kenji, “An Analysis of Requirements Elicitation Process,” IEICE Technical Report, Institute of Electronics, Information and Communication Engineers, vol. 92, No. 128, 47-54 (1992) (with English translation). cited by applicant
Notification of Transmittal of The International Search Report and the Written Opinion of the International Searching Authority, or the Declaration, for Application No. PCT/US09/47735, dated Jul. 23, 2009, 6 pages. cited by applicant
Zamboulis et al. “XML Data Integration by Graph Restructuring,” British National Conference on Databases, Jul. 7, 2004, pp. 57-71. cited by applicant
Office Action in European Appln. No. 17851913.8, mailed Jun. 18, 2024, 18 pages. cited by applicant
Decision of Technical Board of Appeal in European Appln. No. 17851913.8, mailed on Jan. 3, 2025, 25 pages. cited by applicant
Notification of Grant in Singaporean Appln. No. 11201903995V, mailed on Jan. 7, 2025, 3 pages. cited by applicant
Office Action in Japanese Appln No. 2023-216403, mailed on Dec. 11, 2024, 8 pages (with English translation). cited by applicant

Primary Examiner: Richardson; James E

Attorney, Agent or Firm: Fish & Richardson P.C.

Background/Summary

CLAIM OF PRIORITY (1) This application is a continuation application of U.S. patent application Ser. No. 15/829,152, filed on Dec. 1, 2017, which claims priority under 35 U.S.C. § 119 (e) to U.S. Patent Application Ser. No. 62/428,860, filed on Dec. 1, 2016, the entire contents of which are hereby incorporated by reference.

TECHNICAL FIELD

(1) This application relates to data structures and methods for generating, accessing, and displaying lineage metadata, e.g. lineage of an element of data stored in a data storage system.

BACKGROUND

(2) Enterprises use data processing systems, such as data warehousing, customer relationship management, and data mining, to manage data. In many data processing systems, data are pulled from many different data sources, such as database files, operational systems, flat files, the Internet, and other sources into a central repository. Often, data are transformed before being loaded in the data system. Transformation may include cleansing, integration, and extraction. To keep track of data, its sources, and the transformations that have happened to the data stored in a data system, metadata can be used. Metadata (sometimes called “data about data”) are data that describe other data's attributes, format, origins, histories, inter-relationships, etc. Metadata management can play a central role in complex data processing systems.

(3) Sometimes a user may want to investigate how certain data are derived from different data sources. For example, a user may want to know how a dataset or data object was generated or from which source a dataset or data object was imported. Tracing a dataset back to sources from which it is derived is called data lineage tracing (or “upstream data lineage tracing”). Sometimes a user may want to investigate how certain datasets have been used (called “downstream data lineage tracing” or “impact analysis”), for example, which application has read a given dataset. A user may also be

interested in knowing how a dataset is related to other datasets. For example, a user may want to know if a dataset is modified, what tables will be affected.

(4) Lineage, which is a kind of metadata, enables a user to obtain answers to questions about data lineage (e.g., “Where did a given value come from?” “How was the output value computed?” “Which applications produce and depend on this data?”). A user can understand the consequences of proposed modifications (e.g., “If this piece changes, what else will be affected?” “If this source format changes, which applications will be affected?”). A user can also obtain questions to answers involving both technical metadata and business metadata (e.g., “Which groups are responsible for producing and using this data?” “Who changed this application last?” “What changes did they make?”).

(5) The answers to these questions can assist in analyzing and troubleshooting complex data processing systems. For example, if an element of data has an unexpected value, any number of prior inputs or data processing steps might be responsible for this unexpected value. Accordingly, lineage is sometimes presented to a user in the form of a diagram that includes a visual element representing an element of data of interest, as well as visual elements representing other elements of data that affect or are affected by the element of data of interest. A user can view this diagram and visually identify other elements of data and/or transformations that affect the element of data of interest. As an example, by using this information, the user can see whether any of the elements of data and/or transformations may be a source of unexpected values, and correct (or flag for correction) any of the underlying data processing steps if problems are discovered. As another example, by using this information, the user can identify any elements of data or transformations that may be essential to a portion of the system (e.g., such that the element of data of interest would be affected by their removal from the system), and/or elements of data or transformations that may not be essential to a portion of the system (e.g., such that the element of data of interest would not be affected by their removal from the system).

SUMMARY

(6) Among other things, we describe a method performed by a data processing apparatus, the method including receiving a portion of metadata from a data source, the portion of metadata describing nodes and edges, at least some of the edges each representing an effect of one node upon another node, each edge having a single direction; generating instances of a data structure representing the portion of metadata, at least one instance of the data structure including an identification value that identifies a corresponding node, one or more property values representing respective properties of the corresponding node, and one or more pointers to respective identification values, each pointer representing an edge associated with a node identified by the corresponding respective identification value; storing the instances of the data structure in random access memory; receiving a query that includes an identification of at least one particular element of data; and using at least one instance of the data structure to cause a display of a computer system to display a representation of lineage of the particular element of data.

(7) These techniques can be implemented in a number of ways, including as a method, system, and/or computer program product stored on a computer readable storage device.

(8) Aspects of these techniques can include one or more of the following advantages. Lineage metadata can be stored using a special-purpose data structure designed for speed and efficiency when responding to queries for lineage metadata. Lineage metadata can be stored in memory, such that a computer system storing the lineage metadata can respond to queries for lineage metadata more quickly than if the lineage metadata were not stored in memory (e.g., if the lineage metadata were stored in and accessed from a hard disk or another kind of storage technique). Using the techniques described herein, lineage data can be retrieved much faster than other techniques, e.g., 500 times faster.

Description

DESCRIPTION OF DRAWINGS

- (1) FIGS. 1A-1D show a metadata processing environment.
- (2) FIGS. 2A-2E show examples of information displayed in a metadata viewing environment.
- (3) FIG. 3 shows an example data structure.
- (4) FIG. 4 shows an example walk plan.
- (5) FIG. 5A shows a flowchart representing a procedure for storing lineage metadata in a form defined by a special-purpose data structure.
- (6) FIG. 5B shows a flowchart representing a procedure for causing lineage metadata to be displayed.
- (7) FIG. 6 shows a flowchart representing a procedure for traversing lineage metadata stored in the form of special-purpose data structures.
- (8) Like reference symbols in the various drawings indicate like elements.

DESCRIPTION

(9) A system that manages access to metadata can receive a query from a user requesting lineage of a particular element of data and, in response, deliver a diagram representing lineage of the element of data. If the element of data belongs to a data storage system that stores a relatively large amount of data, the system that manages access to the metadata may need to expend a large amount of processing time in order to process the lineage of the element of data and generate the corresponding diagram. However, the processing can be sped up and made more efficient by introducing a system that is dedicated to processing lineage metadata and optimized for this kind of processing. Accordingly, this specification describes a technique by which a specialized system is used for the purpose of processing and storing lineage metadata in a manner that is typically faster and more efficient than if the specialized system were not used.

(10) FIG. 1A shows a metadata processing environment **100** that includes a lineage server **102** which stores and provides lineage metadata to other systems in the environment. The metadata processing environment **100** also includes a metadata server **104** which typically responds to requests for metadata. The metadata server **104** has access to metadata **106** stored in a metadata database **108**. The metadata **106** comes from data sources **110A-C** which contribute metadata **112A-C** to the metadata database **108** on an ongoing basis. For example, the data sources **110A-C** may be any combination of relational databases, flat files, network sources, and so on.

(11) In use, the metadata server **104** responds to a query **114** received from a user terminal **116** operated by a user **118**. For example, the user terminal **116** could be a computing device such as a personal computer, laptop computer, tablet device, smartphone, etc. In some examples, the user terminal **116** operates a network-based user application such as a web browser, e.g., if the metadata server **104** is configured to provide access to data over a network and includes, or communicates with, a web server that can interface with the web browser. In general, many of the interactions between computer systems described herein may take place using the Internet or a similar communications network using communications protocols commonly used on these kinds of networks.

(12) The metadata server **104** is configured to respond to queries for multiple kinds of metadata. Because one kind of metadata is lineage metadata, the metadata server **104** can process a query **114** that requests lineage of a particular element of data, e.g., an element of data **120** stored by one of the data sources **110A-C**, by accessing metadata **106** stored in the metadata database **108** describing the lineage of the particular element of data. The metadata server **104** can then provide lineage metadata **122** to the user terminal **116**, e.g., lineage metadata in the form of a lineage diagram (described below with respect to FIGS. 2A-2E).

(13) In some examples, processing a query **114** related to lineage metadata is a task that takes a

relatively large amount of processing time and/or uses a relatively large amount of processing resources of the metadata server **104**. For example, in order to process the query **114**, the metadata server **104** may need to access metadata **106** stored in the metadata database **108**. In this example, the metadata server **104** would need to spend processing resources generating queries to the metadata database **108** in order to access all of the required metadata. Further, the process of transmitting a query to the metadata database **108** and waiting for a response introduces latency, e.g., communications network latency. Additionally, in some examples, the metadata server **104** would need to process the metadata received from the metadata database **108** in order to extract metadata needed to create a lineage diagram. For example, the metadata received from the metadata database **108** may include information not directly pertinent to lineage of an element of data of interest, since the metadata database **108** stores a variety of kinds of metadata beyond lineage metadata, and so additional processing time is used to identify and remove the information not pertinent to lineage.

(14) In some implementations, the lineage server **102** is used to provide lineage metadata to the metadata server **104**, e.g., to improve performance of the metadata processing environment **100**. The lineage server **102** is a specialized system that stores lineage metadata **124** in a form that can be typically accessed faster and more efficiently than techniques that do not use a lineage server. In particular, the lineage server **102** stores lineage metadata **124** using a special-purpose data structure designed for speed and efficiency when responding to queries for lineage metadata. A data structure defines an arrangement of data, such that all data stored using a particular data structure is arranged in the same manner. The data structure technique is described in further detail below with respect to FIG. 3.

(15) In use, the lineage server **102** transmits queries **126** to the metadata database **108** in order to retrieve lineage metadata **128**. The lineage server **102** ideally stores a comprehensive body of lineage metadata, e.g., for most or all of the data elements **120** stored by the data sources **110A-C**. In this way, the lineage server **102** can respond to queries for lineage for most of the data elements **120** for which a query might be made. As the lineage server **102** receives lineage metadata **128** from the metadata database **108**, the lineage server **102** updates its data structures containing stored lineage metadata **124**. In some examples, the lineage server **102** sends new queries **126** to the metadata database **108** on regular intervals, e.g., every hour or every day or another interval, in order to store a relatively up-to-date body of lineage metadata. For example, the intervals can be scheduled intervals, e.g., corresponding to schedule data maintained by the lineage server **102**.

(16) As shown in greater detail in FIG. 1B, the metadata server **104** receives a query **114** for lineage metadata (e.g., a query for lineage metadata that can be used to display a lineage diagram on the user terminal **116**), the metadata server **104** can provide the query **114** to the lineage server **102**. The lineage server can then return lineage metadata **122** responsive to the query **114**. The metadata server **104** need not expend as much processing time and use as many processing resources preparing the received lineage metadata **122** before providing it to the user terminal **116**, compared to lineage metadata retrieved using other techniques such as retrieving the lineage metadata from the metadata database **108**.

(17) FIGS. 1C and 1D show elements of the lineage server **102** and metadata server **104** and the way in which they interact. As described above, the metadata server **104** receives a query **114** that requests lineage of a particular element of data. The query **114** identifies the particular data element (e.g., one of the data elements **120** of FIG. 1A) for which lineage is requested. The metadata server **104** uses the identity of the data element to select a walk plan **130** from a set of walk plans **132** that can be used to gather lineage metadata related to the data element. A walk plan **130** is a data structure (e.g., a structured document containing tagged portions, such as an XML document) which describes how to traverse (“walk”) a set of lineage metadata in a particular manner. In some examples, a walk plan can be selected based on a data type of the data element. For example, a particular data type can be associated with a particular one of the walk plans **132** (e.g., an

association stored in an index of associations accessible to the metadata server **104**). The walk plans **132** are described in detail below with respect to FIG. 4.

(18) Once a walk plan **130** has been selected, the metadata server **104** transmits the query **114** and the walk plan **130** to the lineage server **102**. In response, the lineage server **102** identifies lineage metadata relevant to the query **114** among its data structures **134** of lineage metadata. The data structures **134** are a representation of lineage metadata arranged in a way that minimizes the amount of storage space needed to contain them, without omitting any data that is needed to respond to a query for lineage metadata. Thus, the lineage server **102** can typically use the data stored in its data structures **134** to provide the metadata server **104** all of the lineage metadata that the metadata server **104** would need to respond to a query **114**. The data structures **134** are described in detail below with respect to FIG. 3.

(19) In some implementations, the data structures **134** are loaded in memory **135** of the lineage server **102** for fast access (e.g., fast reading and writing of data). One example of memory is random access memory. Random access memory stores items of data in a manner such that each item of data can be accessed in substantially the same amount of time as any other item of the same size (e.g., a byte or a word). In contrast, other types of data storage, such as magnetic disks, have physical constraints that cause some items of data to take longer to access than other elements of data, depending on the current physical state of the disk (e.g., the position of a magnetic read/write head). Items of data stored in random access memory are typically stored at an address unique to that item of data or shared among a small number of items of data. Random access memory is typically volatile, such that data stored in random is lost if the random access memory is disconnected from an active power source (e.g., a computer system loses power). In contrast, magnetic disks and some other kinds of data storage is non-volatile and retains data absent an active power source.

(20) Because the lineage server **102** stores the data structures **134** in memory **135**, the lineage server **102** can read and write lineage metadata faster than techniques that do not store the data structures in memory **135**. In particular, the data structures **134** are arranged in a way that minimizes the amount of data use. For example, the data structures **134** may omit data such as text strings present in the original lineage metadata obtained from the metadata server **104**. Thus, all of the data structures **134**, e.g., all of the data representing lineage metadata, can be stored in memory **135** while the lineage server **102** is in use. Computer systems typically have constraints on the amount of random access memory that can be used at a given time (e.g., due to addressing limitations). Further, random access memory tends to be more expensive on a per-byte basis than other types of data storage (e.g., magnetic disk). Thus, if random access memory is used, the data structures **134** may have an upper limit to their combined size on a particular computer system. Accordingly, the techniques described herein (e.g., the techniques described below with respect to FIG. 3) minimize their size but retain information with respect to lineage that may be requested by a query **114**.

(21) In some implementations, the metadata server **104** also stores lineage metadata **137** (e.g., lineage metadata received from the metadata database **108** shown in FIG. 1A). However, the metadata server **104** does not store most of its stored lineage metadata **137** in random access memory, e.g., because the metadata server **104** does not use the data structures of the lineage server **102**. Thus, even if the metadata server **104** also stores some lineage metadata **137**, the metadata server **104** can access the lineage metadata **124** of the lineage server **102** to obtain any metadata not stored locally at the metadata server **104**. If the lineage server **102** were not used, a metadata server **104** storing some lineage metadata **137** typically would access lineage metadata stored in the metadata database **108** (FIG. 1A), which, as described above, may have performance disadvantages compared to using the lineage server **102**.

(22) Although random access memory is used here as the primary example, other types of memory can be used with the lineage server **102**. For example, another kind of memory is flash memory.

Unlike random access memory, flash memory is non-volatile. However, flash memory typically has constraints on accessing items of data. Some types of flash memory are configured in a way that a collection of data items (e.g., blocks of data items) is the smallest unit of data that can be accessed at a time, as opposed to individually accessible data items. For example, in order to delete an item of data on some types of flash memory, an entire block must be deleted. The remaining items of data can be re-written to the flash memory to preserve them.

(23) The lineage server **102** uses the walk plan **130** to traverse the data structures **134** and collect lineage metadata stored in the data structures that is responsive to the query **114**. As shown in FIG. **1D**, the lineage server **102** then sends a response **138** containing the lineage metadata **139** back to the metadata server **104**. The metadata server **104** can use the lineage metadata **139** to generate its own response **140** to the query **114**. The response **140** could take one of several forms. In some examples, the response **140** contains the same lineage metadata **139** received from the lineage server **102**, e.g., in a form with minimal post-processing. In some examples, the metadata server **104** performs post-processing on the lineage metadata **139**. For example, the metadata server **104** may change the form of the lineage metadata **139** to a human-readable form, e.g., if the lineage metadata **139** is received in an encoded format that is not human-readable. In some examples, the metadata server **104** generates a lineage diagram based on the lineage metadata **139** and incorporates data representing the lineage diagram into the response **140**. In some examples, the response **140** is transmitted to the user terminal **116** (FIG. **1A**), e.g., if the response **140** is a lineage diagram (as described in detail below with respect to FIGS. **2A-2E**). In some examples, the response **140** is transmitted to an intermediate system before it is transmitted to the user terminal and/or processed into a form suitable for transmission to the user terminal.

(24) FIG. **2A** shows an example of information displayed in a metadata viewing environment. In some examples, the metadata viewing environment is an interface that executes on a user terminal, e.g., the user terminal **116** shown in FIG. **1A**. In the example of FIG. **2A**, the metadata viewing environment displays information related to a data lineage diagram **200A**. One example of metadata viewing environment is a web-based application that allows a user (e.g., the user **118** shown in FIG. **1A**) to visualize and edit metadata. Using the metadata viewing environment, a user can explore, analyze, and manage metadata using a standard Web browser from anywhere within an enterprise. Each type of metadata object has one or more views or visual representations. The metadata viewing environment of FIG. **2A** illustrates a lineage diagram for target element **206A**.

(25) For example, the lineage diagram displays the end-to-end lineage for the data and/or processing nodes that represent the metadata objects stored in the metadata server **104** (FIG. **1A**); that is, the objects a given starting object depends on (its sources) and the objects that a given starting object affects (its targets). In this example, connections are shown between data elements **202A** and transformations **204A**, two examples of metadata objects. The metadata objects are represented by nodes in the diagram. Data elements **202A** can represent datasets, tables within datasets, columns in tables, and fields in files, messages, and reports, for example. An example of a transformation **204A** is an element of an executable that describes how a single output of a data element is produced. The connections between the nodes are based on relationships among the metadata objects.

(26) FIG. **2B** illustrates a corresponding lineage diagram **200B** for the same target element **206A** shown in FIG. **2A** except each element **202B** is grouped and shown in a group based on a context. For example, data elements **202B** are grouped in datasets **208B** (e.g., tables, files, messages, and reports), applications **210B** (that contain executables such as graphs and plans and programs, plus the datasets that they operate on), and systems **212B**. Systems **212B** are functional groupings of data and the applications that process the data; systems consist of applications and data groups (e.g., databases, file groups, messaging systems, and groups of datasets). Transformations **204B** are grouped in executables **214B**, applications **210B**, and systems **212B**. Executables such as graphs, plans or programs, read and write datasets. Parameters can set what groups are expanded and what

groups are collapsed by default. This allows users to see the details for only the groups that are important to them by removing unnecessary levels of details.

(27) Using the metadata viewing environment to perform data lineage calculations is useful for a number of reasons. For example, calculating and illustrating relationships between data elements and transformations can help a user determine how a reported value was computed for a given field report. A user may also view which datasets store a particular type of data, and which executables read and write to that dataset. In the case of business terms, the data lineage diagram may illustrate which data elements (e.g., columns and fields) are associated with certain business terms (e.g., definitions in an enterprise).

(28) Data lineage diagrams shown within the metadata viewing environment can also aid a user in impact analysis. Specifically, a user may want to know which downstream executables are affected if a column or field is added to a dataset, and who needs to be notified. Impact analysis may determine where a given data element is used, and can also determine the ramifications of changing that data element. Similarly, a user may view what datasets are affected by a change in an executable, or whether it safe to remove a certain database table from production.

(29) Using the metadata viewing environment to perform data lineage calculations for generating data lineage diagrams is useful for business term management. For instance, it is often desirable for employees within an enterprise to agree on the meanings of business terms across that enterprise, the relationships between those terms, and the data to which the terms refer. The consistent use of business terms may enhance the transparency of enterprise data and facilitates communication of business requirements. Thus, it is important to know where the physical data underlying a business term can be found, and what business logic is used in computations.

(30) Viewing relationships between data nodes can also be helpful in managing and maintaining metadata. For instance, a user may wish to know who changed a piece of metadata, what the source (or “source of record”) is for a piece of metadata, or what changes were made when loading or reloading metadata from an external source. In maintaining metadata, it may be desirable to allow designated users to be able to create metadata objects (such as business terms), edit properties of metadata objects (such as descriptions and relationships of objects to other objects), or delete obsolete metadata objects.

(31) The metadata viewing environment provides a number of graphical views of objects, allowing a user to explore and analyze metadata. For example, a user may view the contents of systems and applications and explore the details of any object, and can also view relationships between objects using the data lineage views, which allows a user to easily perform various types of dependency analysis such as the data lineage analysis and impact analysis described above. Hierarchies of objects can also be viewed, and the hierarchies can be searched for specific objects. Once the object is found bookmarks can be created for objects allowing a user to easily return to them.

(32) With the proper permissions, a user can edit the metadata in the metadata viewing environment. For example, a user can update descriptions of objects, create business terms, define relationships between objects (such as linking a business term to a field in a report or column in a table), move objects (for instance, moving a dataset from one application to another) or delete objects.

(33) In FIG. 2C, a corresponding lineage diagram **200C** for target element **206A** is shown, but the level of resolution is set to applications that are participating in the calculation for the target data element **206A**. Specifically, applications **202C**, **204C**, **206C**, **208C**, and **210C** are shown, as only those applications directly participate in the calculation for the target data element **206A**. If a user wishes to view any part of the lineage diagram in a different level of resolution (e.g., to display more or less detail in the diagram), the user may activate the corresponding expand/collapse button **212C**.

(34) FIG. 2D shows a corresponding lineage diagram **200D** at a different level of resolution. In this example, an expand/collapse button **212C** has been activated by a user, and the metadata viewing

environment now displays the same lineage diagram, but application **202C** has been expanded to show the datasets **214D** and executables **216D** within application **202C**.

(35) FIG. **2E** shows a corresponding lineage diagram **200E** at a different level of resolution. In this example, a user has selected to show everything expanded by a custom expansion. Any field or column which is an ultimate source of data (e.g., it has no upstream systems) is expanded. In addition, fields that have a specific flag set are also expanded. In this example, the specific flags are set on datasets and fields at a key intermediate point in the lineage, and one column is the column for which the lineage is being shown.

(36) Other examples of lineage are described in U.S. patent application Ser. No. 12/629,466, titled “VISUALIZING RELATIONSHIPS BETWEEN DATA ELEMENTS AND GRAPHICAL REPRESENTATIONS OF DATA ELEMENT ATTRIBUTES,” which is hereby incorporated by reference in its entirety.

(37) Viewing elements and relationships in the metadata viewing environment can be made more useful by adding information relevant to each of the nodes that represent them. One exemplary way to add relevant information to the nodes is to graphically overlay information on top of certain nodes. These graphics may show some value or characteristic of the data represented by the node, and can be any property in the metadata database. This approach has the advantage of combining two or more normally disparate pieces of information (relationships between nodes of data and characteristics of the data represented by the nodes) and endeavors to put useful information “in context.” For example, characteristics such as metadata quality, metadata freshness, or source of record information can be displayed in conjunction with a visual representation of relationships between data nodes. While some of this information may be accessible in tabular form, it may be more helpful for a user to view characteristics of the data along with the relationships between different nodes of data. A user can select which characteristic of the data will be shown on top of the data element and/or transformation nodes within the metadata viewing environment. Which characteristic is shown can also be set according to default system settings.

(38) As described above with respect to FIG. **1A**, the lineage server **102** uses data structures **134** to store lineage metadata in memory (e.g., random access memory). FIG. **3** shows an example data structure **300**. In use, the lineage server **102** contains many instances of the data structure **300**. An instance of a data structure is a collection of data (e.g., collection of bits) formatted in a manner defined by the data structure. An instance of the data structure **300** described here is sometimes referred to as a “node.”

(39) Each instance of the data structure **300** represents a metadata object, e.g., one of the data elements **202A** or transformations **204A** shown in FIG. **2A**. In some examples, each instance of the data structure **300** represents a node that may be shown in a lineage diagram, e.g., the diagrams **200A-200E** shown in FIGS. **2A-2E**.

(40) In use, the lineage server **102** stores each data structure **300** at a memory location **302** specific to the data structure. Each data structure **300** typically points to memory locations of other data structures.

(41) The data structure **300** is made up of several fields. A field is a collection of data, e.g., a subset of the bits that make up an instance of the data structure **300**. An identifier field **310** includes data representing a unique identifier for an instance of the data structure **300**. A type field **312** includes data representing a type of a metadata object represented by the corresponding instance of the data structure **300**. In some examples, the type could be “data element,” “transformation,” and so on. In some examples, the type field **312** also indicates how many forward and backward edges are included in the instance of the data structure **300**. Properties fields **314** each represent different characteristics of the metadata object represented by the corresponding instance of the data structure **300**. Examples of the properties fields **314** can include a “name” field that includes a text label identifying the metadata object, and a “subtype” field that indicates a subtype of the metadata object, e.g., whether the metadata object represents a file object, executable object, a database

object, or another subtype. Other types of properties can be used. In general, the type field **312** and properties fields **314** can be customized for a particular instance of the lineage server **102**, and are not confined to the examples listed here.

(42) The data structure also includes fields that represent forward edges **316A-C** and backward edges **316D-F**. The edge fields **316A-F** enable the lineage server **102** to “walk” from data structure to data structure and collect the data of the data structure when gathering lineage metadata. In the broadest sense, when we refer to “collecting” a portion data, we mean identifying the portion of data as pertinent to a future action (e.g., transmitting the collected data). Collecting a portion of data sometimes includes copying the data, e.g., copying the data to a buffer or queue to be used in the future action.

(43) Each edge field **316A-F** includes a pointer field **320A-B**. The pointer field **320A-B** stores an address of a respective memory location **322A-B**. In general, a memory location **322A-B** referenced by a pointer field **320A-B** refers to a portion of memory that stores another instance of the data structure **300**. In this way, one instance of a data structure representing a metadata object is “linked” to one or more other instances of data structures representing other metadata objects. Thus, the edges **316A-D** can correspond to, e.g., the relationships among the metadata objects shown in the lineage diagram examples **200A-E** of FIGS. 2A-2E. For example, a forward edge **316A** represents an effect that a metadata object (e.g., the metadata object represented by this instance of the data structure **300**) has on another metadata object (e.g., the metadata object represented by the instance of the data structure at the memory location **322A**). As another example, a backward edge **316D** represents an effect that another metadata object (e.g., the metadata object represented by the instance of the data structure at the memory location **322B**) has on the metadata object of this instance of the data structure **300**.

(44) Each edge field **316A-F** also includes one or more flags **324**. The flags **324** are indicators of information about their associated edge. In some examples, one of the flags **324** may indicate a type of the associated edge, selected from multiple possible types. Many types of edges are possible. For example, some types of edges are input/output edges (representing output from one object and input to another object), element/dataset edges (representing an association between an element and the dataset to which the element belongs), and application/parent edges (representing an association between an executable application and a container, such as a container that also contains datasets associated with the application).

(45) Many of the elements of the data structure **300** typically use a relatively small amount of data. For example, the data associated with the identifier field **310**, type field **312**, and properties fields **314** together may only be a few bytes. e.g., 32 bytes. These fields encode commonly used information within as little as a few bits; for example, if there are only eight possible types for a node, the type field **312** can be as little as three bits long. More complex data, such as strings of text representing the node types, need not be used. Further, the data associated with the memory location **322A-C** is typically the same amount of data as the length of a memory address associated with the type of computer system executing software that instantiates the data structure **300**. Thus, most or all instances of a data structure **300** may use a relatively small amount of data in total, compared to the data used by other techniques for storing lineage metadata.

(46) FIG. 4 shows an example of a walk plan **400**. As described above with respect to FIG. 1C, walk plans **400** are typically stored by the metadata server **104**. In use, the metadata server **104** provides a walk plan to the lineage server **102** when requesting lineage metadata.

(47) A walk plan **400** describes information used by the lineage server **102** when traversing its stored data structures **134**. In general, when a query for lineage metadata is received, e.g., lineage metadata pertinent to a particular metadata object, not all types of lineage metadata need to be returned in response. In some examples, depending on the query, lineage metadata associated with some types of edges may not need to be returned because it is not responsive to the query.

(48) Accordingly, the walk plan **400** includes records **402A-C** for each edge type that may be

among the types of edges represented by the lineage metadata stored by the lineage server **102**. A record **402A** includes an edge type field **404** that includes data indicating the type of edge corresponding to the record **402A**. A record **402A** also includes a follow flag **406**, a collect node flag **408**, and a collect edge flag **409** for the forward direction **410**, and a follow flag **412**, a collect node flag **414**, and a collect edge flag **415** for the backward direction **416**.

(49) A follow flag **406**, **412** indicates whether or not the lineage server **102** should follow an edge of this edge type when traversing its data structures **134**. Put another way, a follow flag **406** for the forward direction **410** indicates whether or not the lineage server **102**, referring to FIG. 3, should access the memory location **322A** identified by a pointer field **320A** of a forward edge field **316A** of an instance of the data structure **300**. Similarly, a follow flag **412** for the backward direction **416** indicates whether or not the lineage server **102**, referring to FIG. 3, should access the memory location **322b** identified by a pointer field **320b** of a backward edge field **316D** of an instance of the data structure **300**.

(50) A collect node flag **408**, **414** indicates whether or not the lineage server **102** should collect an instance of the data structure **300** (FIG. 3), sometimes referred to as a “node,” pointed to by this edge type when traversing its data structures **134**. When we refer to collecting an instance of the data structure **300**, we mean that the data of the instances (or nodes) is added to the data that will be returned in response to a query being processed by the lineage server **102** (FIG. 1A) processing the query. Thus, if a node is collected, data associated with the metadata object represented by the instance of the data structure **300** will be among the lineage metadata returned by the lineage server **102**.

(51) A collect edge flag **409**, **415** indicates whether or not the lineage server **102** should collect the edge (e.g. corresponding to the pointer field **320A** of an instance of the data structure **300**). If an edge is collected, data representing the edge will be among the lineage metadata returned by the lineage server **102**. In some implementations, an edge may not be collected if the edge does not represent a flow of data between the nodes. For example, the edge may represent the association between a data object (represented by one node) and a container of the data object (representing by another node). In this way, by using collect node flags **408**, **414** and collect edge flags **409**, **415** in a walk plan **400**, nodes can be associated with each other in a variety of ways that may or may not be collected for inclusion in lineage metadata, and nodes can represent a variety of data that may or may not be collected for inclusion in lineage metadata.

(52) In some implementations, in use, a walk plan **400** can be represented in the form of one or more XML (Extensible Markup Language) documents. An XML document is a collection of portions separated by “tags.” A tag typically contains a label (e.g., a label identifying the type of tag) and may also include one or more attributes. Tags sometimes come in the form of a start tag and an end tag, such that a start tag is paired with a corresponding end tag. In this way, tags can be hierarchical, such that tags are “nested” within other tags, e.g., by placing a tag between another tag's start tag and end tag pair.

(53) An example walk plan in the form of an XML document is presented below:

(54) TABLE-US-00001 <lineageServerPlan direction=“both”
conditionalOnArg=“!autoFilterEnabled” replacesQueries=“walk”> <useEdge name=“DE-Tr”>
 <condition special=“ExeInterfaceCallStack” /> <condition special=“ControlFilter” />
 <condition special=“Summarization” /> </useEdge> <useEdge name=“Tr-DE”>
<condition special=“ExeInterfaceCallStack” /> <condition special=“ControlFilter” />
<condition special=“Summarization” /> </useEdge> <useEdge name=“DE-DS”
direction=“forward” collectEdge=“false” conditionalOnArg=“walkDSlevel” /> <useEdge
name=“Tr-Exe” direction=“forward” collectEdge=“false” conditionalOnArg=“walkDSlevel” />
<useEdge name=“DS-Exe” conditionalOnArg=“walkDSlevel”> <condition
special=“ExeInterfaceCallStack” /> <condition special=“DSLevelIfNoDE” />
</useEdge> <useEdge name=“Exe-DS” conditionalOnArg=“walkDSlevel”> <condition

```

special="ExeInterfaceCallStack" >      <condition special="DSLevelIfNoDE" />
</useEdge>    <useEdge name="DE-DS" direction="backward" collectEdge="false">
<condition special="DSLevelIfNoDE" />    </useEdge>    <useEdge name="Tr-Exe"
direction="backward" collectEdge="false">      <condition special="DSLevelIfNoDE" >
</useEdge>      </lineageServerPlan>

```

(55) In this example, the “useEdge” tag specifies information for a given type of edge. Each “useEdge” tag can correspond to a record (e.g., the records **402A-C** of the walk plan **400**). The “name” attribute specifies the type of edge (e.g., the edge type **404**), the “direction” attribute specifies the direction (e.g., forward direction **410** or backward direction **416**), the “collectEdge” attribute specifies whether to collect the edge (e.g., the collect flags **408, 414**). Other tags can be used. For example, the “condition special” tag shown in the example above is used to specify custom rules that are carried out when an edge of the specified edge type is followed. In some examples, the custom rules may specify conditions to determine if the edge should be followed and/or collected.

(56) FIG. 5A shows a flowchart representing a procedure **500** for storing lineage metadata in a form defined by a special-purpose data structure, e.g., the data structure **300** shown in FIG. 3. The procedure **500** can be carried out, for example, by components of the lineage server **102** shown in FIG. 1A.

(57) The procedure requests **502** lineage metadata from a metadata source. For example, the metadata source could be the metadata database **108** shown in FIG. 1A. The request could be a request made on regular or semi-regular intervals, for example, every hour, every ten minutes, every minute, or any other interval. In some examples, the request could be made in response to an event, e.g., an event such as a notification that new metadata is available at the metadata source.

(58) Lineage metadata typically describes nodes and edges, such that each node represents a metadata object, and the edges each represent a one-way effect of one node upon another node, e.g., such that each edge has a single direction.

(59) In some examples, e.g., when a lineage server **102** has not yet generated an initial set of data structures representing lineage metadata, the request is a request for all lineage metadata stored by the data source. In some examples, e.g., when the lineage server **102** is updating an existing set of stored data structures, the request is a request for lineage metadata that has been added or changed since the last request.

(60) The procedure receives **504** data, e.g., lineage metadata, from the metadata source. For example, the lineage metadata can be data representing metadata objects and relationships between the metadata objects.

(61) The procedure generates **506** data structures, e.g., instances of the data structure **300** shown in FIG. 3. For example, the data structures can contain information corresponding to the data received from the metadata source. In some examples, each instance of the data structure corresponds to a respective node received from the metadata source. The data structure can include a field for identification values, e.g., an identification value that identifies the node corresponding to an instance of the data structure. The data structure can also include property fields that represent properties of a node corresponding to an instance of the data structure. The data structure can also include pointers to identification values of other nodes, such that the pointers represent edges to the nodes corresponding to the respective identification values.

(62) The procedure stores **508** the data structures. For example, the data structures can be stored in memory, e.g., the memory **135** shown in FIG. 1C. In some examples, the data structures are stored in random access memory. Because the data structures are used to store lineage metadata, any data not relevant to lineage (e.g., other types of metadata stored at the metadata source) can be omitted, reducing the amount of data needed to store the data structures.

(63) In use, the procedure returns to requesting **502** lineage metadata from the metadata source, e.g., on the next regularly scheduled interval.

(64) FIG. 5B shows a flowchart representing a procedure **520** for causing lineage metadata to be displayed. The procedure **520** can be carried out, for example, by components of the lineage server **102** shown in FIG. 1A. In general, a lineage server is configured to return a response to a query that includes metadata describing lineage of a particular element of data, e.g., a metadata object. In some examples, the metadata describes a sequence of nodes and edges, wherein one of the nodes of the sequence represents the particular element of data. In some examples, the procedure **520** is used to access lineage metadata stored by the procedure **500** described above with respect to FIG. 5A.

(65) The procedure receives **522** a query, e.g., a query for lineage metadata. In some examples, the query identifies a metadata object for which lineage metadata is requested.

(66) In some implementations, the query includes an identification of a type of lineage and a walk plan that identifies which types of edges are relevant to the identified type of lineage. In some examples, the walk plan includes conditions for following or collecting an edge based on one or more property values representing respective properties of a corresponding node. An example of a walk plan **400** is shown in FIG. 4.

(67) The procedure gathers **524** lineage metadata. For example, a node representing the metadata object of the received query can be accessed and collected, and edges (e.g., pointers to memory locations) can be traversed to collect other nodes. The gathering of lineage metadata is described in further detail below with respect to FIG. 6.

(68) The procedure transmits **526** the gathered lineage metadata. For example, the gathered lineage metadata can be transmitted to a computer system that issued the query.

(69) After the gathered lineage metadata is transmitted, the gathered lineage metadata may be caused to be displayed **528** on a computer system, e.g., the user terminal **116** shown in FIG. 1A. For example, the lineage metadata may be displayed in the form of a lineage diagram such as the lineage diagrams **200A-200E** shown in FIG. 2A-2E.

(70) FIG. 6 shows a flowchart representing a procedure **600** for traversing lineage metadata stored in the form of special-purpose data structures, e.g., instances of the data structure **300** shown in FIG. 3. The procedure **600** can be carried out, for example, by components of the lineage server **102** shown in FIG. 1A.

(71) The procedure receives **602** a query and walk plan, e.g., the query **114** and walk plan **130** shown in FIG. 1C. The procedure accesses **604** accesses an initial node (e.g., instance of the data structure **300** shown in FIG. 3) representing a metadata object referenced by the query **114**. For example, the initial node may be identified by an identifier field **310** (FIG. 3) storing data that is associated with the metadata object. The initial node is then used as the “current” node, and a recursive portion of the process begins in which the current node is selected from a queue and operations are applied to the current node. Put another way, the initial node is placed in a queue as the first node of the queue, and other nodes are subsequently added to the queue as the procedure is carried out.

(72) The procedure determines **606** if there are remaining forward edge pointers in the current node (e.g., forward edge pointers that have not yet been accessed). If so, the procedure accesses **608** the next pointer that has yet to be accessed, e.g., accesses the memory location of the pointer to retrieve data stored at that memory location. The procedure determines **610** whether to “walk” (e.g., process) the node at that pointer, e.g., according to the edge type associated with the pointer, based on the walk plan (as described above with respect to FIG. 4). If not, the procedure accesses **608** another pointer. If so, the procedure determines whether to collect **611** the node at that pointer. If so, the procedure stores **612** the data of the node to be returned in response to the query, and then places **614** the node in the queue so that its pointers can be accessed. If not, the procedure only puts the node in the queue.

(73) Once all of the forward edge pointers in the current node have been traversed, the procedure determines **616** if there are remaining backward edge pointers in the current node. If so, the procedure accesses **608** the next backward edge pointer.

(74) If there are no remaining forward edge pointers or backward edge pointers, the procedure determines **618** if any nodes remain in the queue. If so, the procedure accesses **620** the next node in the queue, and carries out operations described above using the next node in the queue as the current node. If no nodes remain, the procedure prepares **622** the collected data for transmission to other system. For example, the collected data may be arranged in a particular format because it is transmitted. As another example, encoded data in the collected data may be decoded. For example, data fields containing an encoded value can be converted to a text string corresponding to the value.

(75) Once the data is prepared for transmission, the data can be transmitted, e.g., as described in FIG. 5B with respect to transmission **526** of data.

(76) The systems and techniques described herein can be implemented, for example, using a programmable computing system executing suitable software instructions or it can be implemented in suitable hardware such as a field-programmable gate array (FPGA) or in some hybrid form. For example, in a programmed approach the software may include procedures in one or more computer programs that execute on one or more programmed or programmable computing system (which may be of various architectures such as distributed, client/server, or grid) each including at least one processor, at least one data storage system (including volatile and/or non-volatile memory and/or storage elements), at least one user interface (for receiving input using at least one input device or port, and for providing output using at least one output device or port). The software may include one or more modules of a larger program, for example, that provides services related to the design, configuration, and execution of dataflow graphs. The modules of the program (e.g., elements of a dataflow graph) can be implemented as data structures or other organized data conforming to a data model stored in a data repository.

(77) The software may be stored in non-transitory form, such as being embodied in a volatile or non-volatile storage medium, or any other non-transitory medium, using a physical property of the medium (e.g., surface pits and lands, magnetic domains, or electrical charge) for a period of time (e.g., the time between refresh periods of a dynamic memory device such as a dynamic RAM). In preparation for loading the instructions, the software may be provided on a tangible, non-transitory medium, such as a CD-ROM or other computer-readable medium (e.g., readable by a general or special purpose computing system or device), or may be delivered (e.g., encoded in a propagated signal) over a communication medium of a network to a tangible, non-transitory medium of a computing system where it is executed. Some or all of the processing may be performed on a special purpose computer, or using special-purpose hardware, such as coprocessors or field-programmable gate arrays (FPGAs) or dedicated, application-specific integrated circuits (ASICs). The processing may be implemented in a distributed manner in which different parts of the computation specified by the software are performed by different computing elements. Each such computer program is preferably stored on or downloaded to a computer-readable storage medium (e.g., solid state memory or media, or magnetic or optical media) of a storage device accessible by a general or special purpose programmable computer, for configuring and operating the computer when the storage device medium is read by the computer to perform the processing described herein. The inventive system may also be considered to be implemented as a tangible, non-transitory medium, configured with a computer program, where the medium so configured causes a computer to operate in a specific and predefined manner to perform one or more of the processing steps described herein.

(78) A number of embodiments of the invention have been described. Nevertheless, it is to be understood that the foregoing description is intended to illustrate and not to limit the scope of the invention, which is defined by the scope of the following claims. Accordingly, other embodiments are also within the scope of the following claims. For example, various modifications may be made without departing from the scope of the invention. Additionally, some of the steps described above may be order independent, and thus can be performed in an order different from that described.

Claims

1. A method performed by a data processing system for updating data structures with lineage from a database, wherein updating provides for increased efficiency in determining lineage, including: storing, in memory of a data processing system, a first data structure representing a first item of data or a first transform; wherein the first data structure specifies lineage of a given item of data or a given transform, wherein the lineage specifies an item of data or a transform that affects or is affected by the given item of data or the given transform; and updating, by a data processing system, the lineage specified by the first data structure, by: from time to time, receiving, from a database and by a data processing system, data representing lineage of the given item of data or the given transform; based on the received data representing the lineage, identifying, by a data processing system, a second data structure, stored in memory, representing a second item of data or a second transform that affects or is affected by the given item of data or the given transform; and updating, by a data processing system, the first data structure stored in memory with a reference to the second data structure, stored in memory, representing the second item of data or the second transform that affects or is affected by the given data item or the given transform; wherein the second data structure specifies lineage of the second item of data or the second transform.
2. The method of claim 1, wherein updating the first data structure comprises updating a pointer of the first data structure to reference a memory location of the second data structure representing the second item of data or the second transform.
3. The method of claim 1, further comprising generating a third data structure specifying lineage of the second item of data or the second transform; and storing, in the memory of the data processing system, the third data structure.
4. The method of claim 1, further comprising receiving, from the database, the data representing lineage of the given item of data or the given transform at regular intervals or scheduled intervals.
5. The method of claim 1, wherein the data representing lineage of the given item of data or the given transform is received in response to a request transmitted to the database.
6. The method of claim 1, further comprising: receiving a query including a request for lineage; based on the query, accessing the first data structure and a third data structure to determine the requested lineage; and generating a response to the query, the response including data representing the requested lineage.
7. The method of claim 6, further comprising: accessing a walk plan that includes instructions for determining the requested lineage; and collecting data from the first data structure and third data structure according to the walk plan to determine the requested lineage.
8. The method of claim 1, wherein the given item of data or the given transform is the first item of data or the first transform.
9. A system for updating data structures with lineage from a database, wherein updating provides for increased efficiency in determining lineage, the system comprising: at least one processor; and at least one computer-readable medium storing instructions executable by the at least one processor to perform operations comprising: storing, in memory, a first data structure representing a first item of data or a first transform; wherein the first data structure specifies lineage of a given item of data or a given transform, wherein the lineage specifies an item of data or a transform that affects or is affected by the given item of data or the given transform; and updating the lineage specified by the first data structure, by: from time to time, receiving, from a database, data representing lineage of the given item of data or the given transform; based on the received data representing the lineage, identifying a second data structure, stored in memory, representing a second item of data or a second transform that affects or is affected by the given item of data or the given transform; and updating the first data structure stored in memory with a reference to the second data structure, stored in memory, representing the second item of data or the second transform that affects or is

affected by the given data item or the given transform; wherein the second data structure specifies lineage of the second item of data or the second transform.

10. The system of claim 9, wherein updating the first data structure comprises updating a pointer of the first data structure to reference a memory location of the second data structure representing the second item of data or the second transform.

11. The system of claim 9, wherein the at least one computer readable medium stores instructions executable by the at least one processor to perform operations further comprising: generating a third data structure specifying lineage of the second item of data or the second transform; and storing, in the memory, the third data structure.

12. The system of claim 9, wherein the at least one computer readable medium stores instructions executable by the at least one processor to perform operations further comprising: receiving, from the database, the data representing lineage of the given item of data or the given transform at regular intervals or scheduled intervals.

13. The system of claim 9, wherein the data representing lineage of the given item of data or the given transform is received in response to a request transmitted to the database.

14. The system of claim 9, wherein the at least one computer readable medium stores instructions executable by the at least one processor to perform operations further comprising: receiving a query including a request for lineage; based on the query, accessing the first data structure and a third data structure to determine the requested lineage; and generating a response to the query, the response including data representing the requested lineage.

15. The system of claim 14, wherein the at least one computer readable medium stores instructions executable by the at least one processor to perform operations further comprising: accessing a walk plan that includes instructions for determining the requested lineage; and collecting data from the first data structure and the third data structure according to the walk plan to determine the requested lineage.

16. At least one non-transitory computer readable medium storing instructions executable by at least one processor to perform operations comprising: storing, in memory, a first data structure representing a first item of data or a first transform; wherein the first data structure specifies lineage of a given item of data or a given transform, wherein the lineage specifies an item of data or a transform that affects or is affected by the given item of data or the given transform; and updating the lineage specified by the first data structure, by: from time to time, receiving, from a database, data representing lineage of the given item of data or the given transform; based on the received data representing the lineage, identifying a second data structure, stored in memory, representing a second item of data or a second transform that affects or is affected by the given item of data or the given transform; and updating the first data structure stored in memory with a reference to the second data structure, stored in memory, representing the second item of data or the second transform that affects or is affected by the given data item or the given transform; wherein the second data structure specifies lineage of the second item of data or the second transform.

17. The at least one non-transitory computer readable medium of claim 16, wherein updating the first data structure comprises updating a pointer of the first data structure to reference a memory location of the second data structure representing the second item of data or the second transform.

18. The at least one non-transitory computer readable medium of claim 16 storing instructions executable by the at least one processor to perform operations further comprising: generating a third data structure specifying lineage of the second item of data or the second transform; and storing, in the memory, the third data structure.

19. The at least one non-transitory computer readable medium of claim 16 storing instructions executable by the at least one processor to perform operations further comprising: receiving, from the database, the data representing lineage of the given item of data or the given transform at regular intervals or scheduled intervals.

20. The at least one non-transitory computer readable medium of claim 16, wherein the data

representing lineage of the given item of data or the given transform is received in response to a request transmitted to the database.

21. The at least one non-transitory computer readable medium of claim 16 storing instructions executable by the at least one processor to perform operations further comprising: receiving a query including a request for lineage; based on the query, accessing the first data structure and a third data structure to determine the requested lineage; and generating a response to the query, the response including data representing the requested lineage.
