

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405955
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Gong; Qing et al.

Automatically transforming data into a compacted electronic data structure

Abstract

One example described herein can involve a system that receives, from a user, a structured document that describes one or more conditions selected by the user to customize one or more electronic data structures. The system can also query a database to retrieve data from the database. The system can then execute an interpreter program, which can ingest the structured document and the data, identify a subset of the data that satisfies the condition(s), and generate the one or more electronic data structures based on the subset of the data. The electronic data structure(s) can include the identified subset of the data and exclude a remainder of the data. The system can compact the electronic data structure(s) by removing empty fields therein. The system can then provide the compacted data structure(s) to the user.

Inventors: Gong; Qing (Cary, NC), Brizzotti; Murilo Machado (Cary, NC), Lewis; Kimberly W. (Morehead City, NC), Applegate; David C. (Sparta, NC), Peterson; Chad Delano (Cary, NC)

Applicant: SAS Institute Inc. (Cary, NC)

Family ID: 96587162

Assignee: SAS INSTITUTE, INC. (Cary, NC)

Appl. No.: 18/751452

Filed: June 24, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250252104 A1	Aug. 07, 2025

Related U.S. Application Data

Publication Classification

Int. Cl.: G06F16/2455 (20190101); G06F16/248 (20190101); G06F16/81 (20190101)

U.S. Cl.:

CPC G06F16/2455 (20190101); G06F16/248 (20190101); G06F16/81 (20190101);

Field of Classification Search

CPC: G06F (16/2455); G06F (16/248); G06F (16/81)

USPC: 707/769

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7743071	12/2009	Yang et al.	N/A	N/A
9128908	12/2014	Weller	N/A	N/A
9223847	12/2014	Aski et al.	N/A	N/A
9450604	12/2015	Kumaran	N/A	H03M 7/3082
2002/0123922	12/2001	Poage	705/7.29	G06Q 30/0201
2006/0174192	12/2005	Mortensen	N/A	N/A
2010/0070491	12/2009	Cragun	707/722	G06F 16/00
2010/0325169	12/2009	Loh	707/812	G06F 16/81
2013/0290829	12/2012	Peng et al.	N/A	N/A
2020/0364269	12/2019	Kwan	N/A	G06F 16/1847

OTHER PUBLICATIONS

Popoola, et al. “Code Synthesis for Sparse Tensor Format Conversion and Optimization” Conference 2017, Jul. 2017, Washington, DC, USA. cited by applicant

Sanderson, et al. “Practical Sparse Matrices in C++ with Hybrid Storage and Template-Based Expression Optimisation” Practical Sparse Matrices in C++ with Hybrid Storage and Template-Based Expression Optimisation. Mathematical and Computational Applications, vol. 24, No. 3, 2019. <https://doi.org/10.3390/mca24030070>. cited by applicant

“Convert Sparse Matrix to Full Storage—MATLAB Full” Mar. 18, 2024, 4 pages, <https://www.mathworks.com/help/matlab/ref/full.html>. cited by applicant

“Creating Compact View of a Sparse Table in Tableau” Mar. 18, 2024, 5 pages, <https://vizdiff.blogspot.com/2022/01/compacting-sparse-table-in-tableau.html>. cited by applicant

Fredman, et al. “Storing a Sparse Table with O(1) Worst Case Access Time” Journal of the Association for Computing Machinery, vol. 31, No. 3, Jul. 1984, pp. 538-544. cited by applicant

“Eigen: Sparse Matrix Manipulation” Mar. 18, 2024; 8 pages, https://eigen.tuxfamily.org/dox/group_TutorialSparse.html. cited by applicant

“Fast Character Case Conversion, or how to Really Compress Sparse Arrays” Mar. 18, 2024, 3 pages, <https://news.ycombinator.com/item?id=26227764>. cited by applicant

Cheshmi, et al. “ParSy: Inspection and Transformation of Sparse Matrix Computations for

Parallelism” Jun. 11, 2024, 13 pages. cited by applicant

Chou, et al. “Automatic Generation of Efficient Sparse Tensor Format Conversion Routines” In Proceedings of the 41st ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI '20), Jun. 15-20, 2020, London, UK. ACM, New York, NY, USA, 16 pages. <https://doi.org/10.1145/3385412.3385963>. cited by applicant

Fujiki, et al. “Near-Memory Data Transformation for Efficient Sparse Matrix Multi-Vector Multiplication” SC '19, Nov. 17-22, 2019, Denver, CO, USA. cited by applicant

Zhao, et al. “Processing and Storing Sparse Data in SAS® Using SAS® Text Miner Procedures” SAS Institute, Inc. Paper SAS195-2014, 13 pages. cited by applicant

King, et al. “Dynamic Sparse-Matrix Allocation on GPUs” University of Utah, Jun. 11, 2024, 19 pages. cited by applicant

Tobi Goodness Popoola—“Sparse Format Conversion and Code Synthesis” A dissertation submitted in partial fulfillment of the requirements for the degree of Doctor of Philosophy in Computing, Boise State University May 2023, 107 pages. cited by applicant

Primary Examiner: Alam; Shahid A

Attorney, Agent or Firm: Kilpatrick Townsend & Stockton, LLP

Background/Summary

REFERENCE TO RELATED APPLICATION (1) This claims the benefit of priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 63/548,780 filed Feb. 1, 2024, the entirety of which is hereby incorporated by reference herein.

Technical Field

(1) The present disclosure relates generally to using a computer to automatically transform data. More specifically, but not by way of limitation, this disclosure relates to a computerized approach for automatically transforming data into an electronic data structure that can be customized using a structured document and compacting the electronic data structure to reduce storage (e.g., memory and disk storage) consumption.

Background

(2) There can be different types of data stored in a database. One common type of data that can be stored in a database is a sparse dataset. A sparse dataset can be a dataset in which a high percentage of the data is missing or is set to a value of zero. In some cases, it may be desirable to transform the data from its current format in the database to a new format. For example, a computer program may only accept inputs in a particular format, which may be different from the format in which the data is stored in the database. Providing the data to the computer program in the wrong format may lead to errors and/or failures. Thus, before the data is input into the computer program, the data may need to be converted from its existing format in the database to the particular format that is compatible with the computer program. Different computer programs may require the data to be formatted in different ways, some or all of which may be different from how the data is stored in the database.

SUMMARY

(3) One example of the present disclosure includes a method of operations implemented by one or more processors. The operations can include receiving a structured document from a user, wherein the structured document describes one or more conditions selected by the user to customize one or more electronic data structures, and wherein the one or more conditions are defined in the structured document using a declarative language. The operations can include querying a database

to retrieve data from the database. The operations can include ingesting, by an interpreter program, the structured document and the data. The operations can include identifying, by the interpreter program using the structured document and the data, a subset of the data that satisfies the one or more conditions. The operations can include generating, by the interpreter program, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data. The operations can include compacting the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures. The operations can include providing the one or more compacted data structures to the user.

(4) Another example of the present disclosure can include a system comprising one or more processors and one or more memories. The one or more memories can include program code that is executable by the one or more processors for causing the one or more processors to perform operations. The operations can include receiving a structured document from a user, wherein the structured document describes one or more conditions selected by the user to customize one or more electronic data structures, and wherein the one or more conditions are defined in the structured document using a declarative language. The operations can include querying a database to retrieve data from the database. The operations can include ingesting, by an interpreter program, the structured document and the data. The operations can include identifying, by the interpreter program using the structured document and the data, a subset of the data that satisfies the one or more conditions. The operations can include generating, by the interpreter program, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data. The operations can include compacting the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures. The operations can include providing the one or more compacted data structures to the user.

(5) Yet another example of the present disclosure can include a non-transitory computer-readable medium comprising program code that is executable by the one or more processors for causing the one or more processors to perform operations. The operations can include receiving a structured document from a user, wherein the structured document describes one or more conditions selected by the user to customize one or more electronic data structures, and wherein the one or more conditions are defined in the structured document using a declarative language. The operations can include querying a database to retrieve data from the database. The operations can include ingesting, by an interpreter program, the structured document and the data. The operations can include identifying, by the interpreter program using the structured document and the data, a subset of the data that satisfies the one or more conditions. The operations can include generating, by the interpreter program, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data. The operations can include compacting the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures. The operations can include providing the one or more compacted data structures to the user.

(6) This summary is not intended to identify key or essential features of the claimed subject matter, nor is it intended to be used in isolation to determine the scope of the claimed subject matter. The subject matter should be understood by reference to appropriate portions of the entire specification, any or all drawings, and each claim.

(7) The foregoing, together with other features and examples, will become more apparent upon referring to the following specification, claims, and accompanying drawings.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure is described in conjunction with the appended figures:
- (2) FIG. 1 shows an example of the hardware components of a data transmission network according to some aspects of the present disclosure.
- (3) FIG. 2 shows an example network including an example set of devices communicating with each other over an exchange system according to some aspects of the present disclosure.
- (4) FIG. 3 shows an example representation of a conceptual model of a communications protocol system according to some aspects of the present disclosure.
- (5) FIG. 4 shows a communications grid computing system including a variety of control and worker nodes according to some aspects of the present disclosure.
- (6) FIG. 5 shows a flow chart showing an example process for adjusting a communications grid or a work project in a communications grid after a failure of a node according to some aspects of the present disclosure.
- (7) FIG. 6 shows a portion of a communications grid computing system including a control node and a worker node according to some aspects of the present disclosure.
- (8) FIG. 7 shows a flow chart showing an example method **700** for executing a project within a grid computing system according to some aspects of the present disclosure.
- (9) FIG. 8 shows a block diagram including components of an Event Stream Processing Engine (ESPE) according to some aspects of the present disclosure.
- (10) FIG. 9 shows a flow chart of an example process including operations performed by an event stream processing engine according to some aspects of the present disclosure.
- (11) FIG. 10 shows an ESP system interfacing between publishing device and event subscribing devices according to some aspects of the present disclosure.
- (12) FIG. 11 shows a flow chart of an example of a process for generating and using a machine-learning model according to some aspects of the present disclosure.
- (13) FIG. 12 shows a node-link diagram of an example of a neural network according to some aspects of the present disclosure.
- (14) FIG. 13 shows various aspects of the use of containers as a mechanism to allocate processing, storage and/or other resources of a processing system to the performance of various analyses according to some aspects of the present disclosure.
- (15) FIG. 14 shows a block diagram of an example of a system for automatically transforming data into an electronic data structure according to some aspects of the present disclosure.
- (16) FIG. 15 shows a flowchart of an example of a process for automatically transforming data into an electronic data structure according to some aspects of the present disclosure.
- (17) FIG. 16 shows an example of data that can be transformed into a custom electronic data structure according to some aspects of the present disclosure.
- (18) FIG. 17 shows an example of a structured document usable to transform data into a custom electronic data structure according to some aspects of the present disclosure.
- (19) FIG. 18 shows an example of a structured document usable to transform data into a custom electronic data structure according to some aspects of the present disclosure.
- (20) FIG. 19 shows an example of transforming data into multiple custom electronic data structures according to some aspects of the present disclosure.
- (21) FIG. 20 shows an example of compacting an electronic data structure according to some aspects of the present disclosure.
- (22) In the appended figures, similar components or features can have the same reference number. Further, various components of the same type may be distinguished by following the reference number with a lowercase letter that distinguishes among the similar components. If only the first

reference number is used in the specification, the description is applicable to any one of the similar components having the same first reference number irrespective of the lowercase letter.

DETAILED DESCRIPTION

(23) Certain aspects and features of the present disclosure relate to a computerized approach for automatically transforming data into an electronic data structure that can be customized using a structured document, and compacting the electronic data structure to reduce storage consumption and improve usability. For example, a user can customize how they want the data to be presented in the electronic data structure by including one or more conditions in the structured document. The structured document can be written in a declarative format to make it easier for a human user to define their customizations. For instance, the structured document can be an eXtensible Markup Language (XML) or JavaScript Object Notation (JSON) file in which the user has input their one or more customization conditions. A computer system can receive the structured document and obtain the data, for example by querying a database. The computer system can then execute an interpreter program configured to ingest the data and the structured document. Using these inputs, the interpreter program can identify a subset of the data that satisfies the conditions defined in the structured document. The interpreter program can then generate one or more electronic data structures that include that identified subset of the data and exclude a remainder of the data. In this way, the interpreter program can automatically generate the one or more electronic data structures in accordance with the preferences of the user as defined in the structured document. In some examples, the computer system can next perform a compacting process that removes one or more empty fields from the one or more electronic data structures. An empty field can be a field with no value. With the empty fields removed, the compacted data structures can consume less memory than the original electronic data structures. The system can then present the one or more compacted data structures to the user.

(24) In some examples, the interpreter program can generate two or more electronic data structures based on the subset of the data. Such electronic data structures may be the same as or different from one another in format, content, and/or structure. For instance, based on the content of the structured document, the interpreter program can generate three electronic data structures that differ from one another in format, content, and structure. Thus, a single interpreter program may be capable of dividing a single dataset into multiple electronic data structures with different characteristics from one another.

(25) The structured document described above can provide a relatively simple and flexible way for a user to customize which data is included in the electronic data structure. This may allow a greater number of users to be able to customize their electronic data structures, because no knowledge of computer programming is required. In contrast, conventional approaches may involve an interpreter program that has its settings hardcoded in the source code, including which data to include in the electronic data structure. This is a rigid approach that requires knowledge of computer programming to modify the source code and recompile the interpreter program, if the user wants to adjust which data is to be included in the electronic data structure.

(26) Additionally, some kinds of data may have a large number of empty fields. One example of such data can be a sparse dataset, which in some cases can have 30%, 50%, 70%, or more of its fields be empty. Maintaining those empty fields in the electronic data structures may unnecessarily consume memory, since those empty fields may not have any valuable data and may effectively be ignored by many software programs (e.g., machine-learning models) that use the electronic data structures. Accordingly, some examples described herein can perform a compaction process that removes some or all of the empty fields to reduce unnecessary storage consumption. For instance, the compaction process may remove an entire column or empty fields, an entire row of empty fields, or both. This may help conserve the limited memory space of the computer system for more important data.

(27) In some examples, there can be multiple different interpreter programs that may each produce

a different type of electronic data structure. For instance, each interpreter program may produce an electronic data structure that is in a different format from the other interpreter programs. In some such examples, the computer system can determine a target format into which the data is to be organized in the electronic data structure. The computer system can make this determination based on the structured document, which may define the target format for the electronic data structure. Based on this determination, the computer system can select a particular interpreter program, from among the multiple available interpreter programs, that is configured to generate an electronic data structure in the target format. The computer system can then use the selected interpreter program to generate the electronic data structure in the target format.

(28) Additionally, or alternatively, the different interpreter programs may be configured to operate on different types of input data. For instance, each interpreter program may be configured to operate on input data that is in certain format. In some such examples, the computer system can determine a format of the data and select a particular interpreter program, from among the multiple available interpreter programs, that is configured to operate on data in that format. The computer system can then use the selected interpreter program to generate the electronic data structure based on the data. In this way, the computer system may account for the format of the input data, the format for the output electronic data structure, or both in selecting an appropriate interpreter program. This may help ensure compatibility between the input data and the interpreter program. This may also help ensure that the output electronic data structure satisfies the user's objectives.

(29) These illustrative examples are given to introduce the reader to the general subject matter discussed here and are not intended to limit the scope of the disclosed concepts. The following sections describe various additional features and examples with reference to the drawings in which like numerals indicate like elements but, like the illustrative examples, should not be used to limit the present disclosure.

(30) FIG. 1 is a block diagram that provides an illustration of the hardware components of a data transmission network **100**, according to embodiments of the present technology. Data transmission network **100** is a specialized computer system that may be used for processing large amounts of data where a large number of computer processing cycles are required.

(31) Data transmission network **100** may also include computing environment **114**. Computing environment **114** may be a specialized computer or other machine that processes the data received within the data transmission network **100**. Data transmission network **100** also includes one or more network devices **102**. Network devices **102** may include client devices that attempt to communicate with computing environment **114**. For example, network devices **102** may send data to the computing environment **114** to be processed, may send signals to the computing environment **114** to control different aspects of the computing environment or the data it is processing, among other reasons. Network devices **102** may interact with the computing environment **114** through a number of ways, such as, for example, over one or more networks **108**. As shown in FIG. 1, computing environment **114** may include one or more other systems. For example, computing environment **114** may include a database system **118** and/or a communications grid **120**.

(32) In other embodiments, network devices **102** may provide a large amount of data, either all at once or streaming over a period of time (e.g., using event stream processing (ESP), described further with respect to FIGS. 8-10), to the computing environment **114** via networks **108**. For example, network devices **102** may include network computers, sensors, databases, or other devices that may transmit or otherwise provide data to computing environment **114**. For example, network devices **102** may include local area network devices, such as routers, hubs, switches, or other computer networking devices. These devices may provide a variety of stored or generated data, such as network data or data specific to the network devices themselves. Network devices **102** may also include sensors that monitor their environment or other devices to collect data regarding that environment or those devices, and such network devices may provide data they collect overtime. Network devices **102** may also include devices within the internet of things, such as devices within

a home automation network. Some of these devices may be referred to as edge devices, and may involve edge computing circuitry. Data may be transmitted by network devices **102** directly to computing environment **114** or to network-attached data stores, such as network-attached data stores **110** for storage so that the data may be retrieved later by the computing environment **114** or other portions of data transmission network **100**.

(33) Data transmission network **100** may also include one or more network-attached data stores **110**. Network-attached data stores **110** are used to store data to be processed by the computing environment **114** as well as any intermediate or final data generated by the computing system in non-volatile memory. However, in certain embodiments, the configuration of the computing environment **114** allows its operations to be performed such that intermediate and final data results can be stored solely in volatile memory (e.g., RAM), without a requirement that intermediate or final data results be stored to non-volatile types of memory (e.g., disk). This can be useful in certain situations, such as when the computing environment **114** receives ad hoc queries from a user and when responses, which are generated by processing large amounts of data, need to be generated on-the-fly. In this non-limiting situation, the computing environment **114** may be configured to retain the processed information within memory so that responses can be generated for the user at different levels of detail as well as allow a user to interactively query against this information.

(34) Network-attached data stores **110** may store a variety of different types of data organized in a variety of different ways and from a variety of different sources. For example, network-attached data storage may include storage other than primary storage located within computing environment **114** that is directly accessible by processors located therein. Network-attached data storage may include secondary, tertiary or auxiliary storage, such as large hard drives, servers, virtual memory, among other types. Storage devices may include portable or non-portable storage devices, optical storage devices, and various other mediums capable of storing or containing data. A machine-readable storage medium or computer-readable storage medium may include a non-transitory medium in which data can be stored and that does not include carrier waves and/or transitory electronic signals. Examples of a non-transitory medium may include, for example, a magnetic disk or tape, optical storage media such as compact disk or digital versatile disk, flash memory, memory or memory devices. A computer-program product may include code and/or machine-executable instructions that may represent a procedure, a function, a subprogram, a program, a routine, a subroutine, a module, a software package, a class, or any combination of instructions, data structures, or program statements. A code segment may be coupled to another code segment or a hardware circuit by passing and/or receiving information, data, arguments, parameters, or memory contents. Information, arguments, parameters, data, etc. may be passed, forwarded, or transmitted via any suitable means including memory sharing, message passing, token passing, and network transmission, among others. Furthermore, the data stores may hold a variety of different types of data. For example, network-attached data stores **110** may hold unstructured (e.g., raw) data, such as manufacturing data (e.g., a database containing records identifying products being manufactured with parameter data for each product, such as colors and models) or product sales databases (e.g., a database containing individual data records identifying details of individual product sales).

(35) The unstructured data may be presented to the computing environment **114** in different forms such as a flat file or a conglomerate of data records, and may have data values and accompanying time stamps. The computing environment **114** may be used to analyze the unstructured data in a variety of ways to determine the best way to structure (e.g., hierarchically) that data, such that the structured data is tailored to a type of further analysis that a user wishes to perform on the data. For example, after being processed, the unstructured time stamped data may be aggregated by time (e.g., into daily time period units) to generate time series data and/or structured hierarchically according to one or more dimensions (e.g., parameters, attributes, and/or variables). For example, data may be stored in a hierarchical data structure, such as a ROLAP OR MOLAP database, or may

be stored in another tabular form, such as in a flat-hierarchy form.

(36) Data transmission network **100** may also include one or more server farms **106**. Computing environment **114** may route select communications or data to the one or more server farms **106** or one or more servers within the server farms. Server farms **106** can be configured to provide information in a predetermined manner. For example, server farms **106** may access data to transmit in response to a communication. Server farms **106** may be separately housed from each other device within data transmission network **100**, such as computing environment **114**, and/or may be part of a device or system.

(37) Server farms **106** may host a variety of different types of data processing as part of data transmission network **100**. Server farms **106** may receive a variety of different data from network devices **102**, from computing environment **114**, from cloud network **116**, or from other sources. The data may have been obtained or collected from one or more sensors, as inputs from a control database, or may have been received as inputs from an external system or device. Server farms **106** may assist in processing the data by turning raw data into processed data based on one or more rules implemented by the server farms. For example, sensor data may be analyzed to determine changes in an environment over time or in real-time.

(38) Data transmission network **100** may also include one or more cloud networks **116**. Cloud network **116** may include a cloud infrastructure system that provides cloud services. In certain embodiments, services provided by the cloud network **116** may include a host of services that are made available to users of the cloud infrastructure system on demand. Cloud network **116** is shown in FIG. 1 as being connected to computing environment **114** (and therefore having computing environment **114** as its client or user), but cloud network **116** may be connected to or utilized by any of the devices in FIG. 1. Services provided by the cloud network can dynamically scale to meet the needs of its users. The cloud network **116** may include one or more computers, servers, and/or systems. In some embodiments, the computers, servers, and/or systems that make up the cloud network **116** are different from the user's own on-premises computers, servers, and/or systems. For example, the cloud network **116** may host an application, and a user may, via a communication network such as the Internet, on demand, order and use the application.

(39) While each device, server and system in FIG. 1 is shown as a single device, it will be appreciated that multiple devices may instead be used. For example, a set of network devices can be used to transmit various communications from a single user, or remote server may include a server stack. As another example, data may be processed as part of computing environment **114**.

(40) Each communication within data transmission network **100** (e.g., between client devices, between servers **106** and computing environment **114** or between a server and a device) may occur over one or more networks **108**. Networks **108** may include one or more of a variety of different types of networks, including a wireless network, a wired network, or a combination of a wired and wireless network. Examples of suitable networks include the Internet, a personal area network, a local area network (LAN), a wide area network (WAN), or a wireless local area network (WLAN). A wireless network may include a wireless interface or combination of wireless interfaces. As an example, a network in the one or more networks **108** may include a short-range communication channel, such as a BLUETOOTH® communication channel or a BLUETOOTH® Low Energy communication channel. A wired network may include a wired interface. The wired and/or wireless networks may be implemented using routers, access points, bridges, gateways, or the like, to connect devices in the network **108**, as will be further described with respect to FIG. 2. The one or more networks **108** can be incorporated entirely within or can include an intranet, an extranet, or a combination thereof. In one embodiment, communications between two or more systems and/or devices can be achieved by a secure communications protocol, such as secure sockets layer (SSL) or transport layer security (TLS). In addition, data and/or transactional details may be encrypted.

(41) Some aspects may utilize the Internet of Things (IoT), where things (e.g., machines, devices, phones, sensors) can be connected to networks and the data from these things can be collected and

processed within the things and/or external to the things. For example, the IoT can include sensors in many different devices, and high value analytics can be applied to identify hidden relationships and drive increased efficiencies. This can apply to both big data analytics and real-time (e.g., ESP) analytics. This will be described further below with respect to FIG. 2.

(42) As noted, computing environment **114** may include a communications grid **120** and a transmission network database system **118**. Communications grid **120** may be a grid-based computing system for processing large amounts of data. The transmission network database system **118** may be for managing, storing, and retrieving large amounts of data that are distributed to and stored in the one or more network-attached data stores **110** or other data stores that reside at different locations within the transmission network database system **118**. The compute nodes in the grid-based computing system **120** and the transmission network database system **118** may share the same processor hardware, such as processors that are located within computing environment **114**.

(43) FIG. 2 illustrates an example network including an example set of devices communicating with each other over an exchange system and via a network, according to embodiments of the present technology. As noted, each communication within data transmission network **100** may occur over one or more networks. System **200** includes a network device **204** configured to communicate with a variety of types of client devices, for example client devices **230**, over a variety of types of communication channels.

(44) As shown in FIG. 2, network device **204** can transmit a communication over a network (e.g., a cellular network via a base station). The communication can be routed to another network device, such as network devices **205-209**, via base station. The communication can also be routed to computing environment **214** via base station. For example, network device **204** may collect data either from its surrounding environment or from other network devices (such as network devices **205-209**) and transmit that data to computing environment **214**.

(45) Although network devices **204-209** are shown in FIG. 2 as a mobile phone, laptop computer, tablet computer, temperature sensor, motion sensor, and audio sensor respectively, the network devices may be or include sensors that are sensitive to detecting characteristics of their environment. For example, the network devices may include sensors such as water sensors, power sensors, electrical current sensors, chemical sensors, optical sensors, pressure sensors, geographic or position sensors (e.g., GPS), velocity sensors, acceleration sensors, flow rate sensors, among others. Examples of characteristics that may be sensed include force, torque, load, strain, position, temperature, air pressure, fluid flow, chemical properties, resistance, electromagnetic fields, radiation, irradiance, proximity, acoustics, moisture, distance, speed, vibrations, acceleration, electrical potential, and electrical current, among others. The sensors may be mounted to various components used as part of a variety of different types of systems (e.g., an oil drilling operation). The network devices may detect and record data related to the environment that it monitors, and transmit that data to computing environment **214**.

(46) As noted, one type of system that may include various sensors that collect data to be processed and/or transmitted to a computing environment according to certain embodiments includes an oil drilling system. For example, the one or more drilling operation sensors may include surface sensors that measure a hook load, a fluid rate, a temperature and a density in and out of the wellbore, a standpipe pressure, a surface torque, a rotation speed of a drill pipe, a rate of penetration, a mechanical specific energy, etc., and downhole sensors that measure a rotation speed of a bit, fluid densities, downhole torque, downhole vibration (axial, tangential, lateral), a weight applied at a drill bit, an annular pressure, a differential pressure, an azimuth, an inclination, a dog leg severity, a measured depth, a vertical depth, a downhole temperature, etc. Besides the raw data collected directly by the sensors, other data may include parameters either developed by the sensors or assigned to the system by a client or other controlling device. For example, one or more drilling operation control parameters may control settings such as a mud motor speed to flow ratio, a bit diameter, a predicted formation top, seismic data, weather data, etc. Other data may be generated

using physical models such as an earth model, a weather model, a seismic model, a bottom hole assembly model, a well plan model, an annular friction model, etc. In addition to sensor and control settings, predicted outputs, of for example, the rate of penetration, mechanical specific energy, hook load, flow in fluid rate, flow out fluid rate, pump pressure, surface torque, rotation speed of the drill pipe, annular pressure, annular friction pressure, annular temperature, equivalent circulating density, etc. may also be stored in the data warehouse.

(47) In another example, another type of system that may include various sensors that collect data to be processed and/or transmitted to a computing environment according to certain embodiments includes a home automation or similar automated network in a different environment, such as an office space, school, public space, sports venue, or a variety of other locations. Network devices in such an automated network may include network devices that allow a user to access, control, and/or configure various home appliances located within the user's home (e.g., a television, radio, light, fan, humidifier, sensor, microwave, iron, and/or the like), or outside of the user's home (e.g., exterior motion sensors, exterior lighting, garage door openers, sprinkler systems, or the like). For example, network device **102** may include a home automation switch that may be coupled with a home appliance. In another embodiment, a network device can allow a user to access, control, and/or configure devices, such as office-related devices (e.g., copy machine, printer, or fax machine), audio and/or video related devices (e.g., a receiver, a speaker, a projector, a DVD player, or a television), media-playback devices (e.g., a compact disc player, a CD player, or the like), computing devices (e.g., a home computer, a laptop computer, a tablet, a personal digital assistant (PDA), a computing device, or a wearable device), lighting devices (e.g., a lamp or recessed lighting), devices associated with a security system, devices associated with an alarm system, devices that can be operated in an automobile (e.g., radio devices, navigation devices), and/or the like. Data may be collected from such various sensors in raw form, or data may be processed by the sensors to create parameters or other data either developed by the sensors based on the raw data or assigned to the system by a client or other controlling device.

(48) In another example, another type of system that may include various sensors that collect data to be processed and/or transmitted to a computing environment according to certain embodiments includes a power or energy grid. A variety of different network devices may be included in an energy grid, such as various devices within one or more power plants, energy farms (e.g., wind farm, solar farm, among others) energy storage facilities, factories, homes and businesses of consumers, among others. One or more of such devices may include one or more sensors that detect energy gain or loss, electrical input or output or loss, and a variety of other efficiencies. These sensors may collect data to inform users of how the energy grid, and individual devices within the grid, may be functioning and how they may be made more efficient.

(49) Network device sensors may also perform processing on data they collect before transmitting the data to the computing environment **114**, or before deciding whether to transmit data to the computing environment **114**. For example, network devices may determine whether data collected meets certain rules, for example by comparing data or values calculated from the data and comparing that data to one or more thresholds. The network device may use this data and/or comparisons to determine if the data should be transmitted to the computing environment **214** for further use or processing.

(50) Computing environment **214** may include machines **220** and **240**. Although computing environment **214** is shown in FIG. 2 as having two machines, **220** and **240**, computing environment **214** may have only one machine or may have more than two machines. The machines that make up computing environment **214** may include specialized computers, servers, or other machines that are configured to individually and/or collectively process large amounts of data. The computing environment **214** may also include storage devices that include one or more databases of structured data, such as data organized in one or more hierarchies, or unstructured data. The databases may communicate with the processing devices within computing environment **214** to distribute data to

them. Since network devices may transmit data to computing environment **214**, that data may be received by the computing environment **214** and subsequently stored within those storage devices. Data used by computing environment **214** may also be stored in data stores **235**, which may also be a part of or connected to computing environment **214**.

(51) Computing environment **214** can communicate with various devices via one or more routers **225** or other inter-network or intra-network connection components. For example, computing environment **214** may communicate with client devices **230** via one or more routers **225**.

Computing environment **214** may collect, analyze and/or store data from or pertaining to communications, client device operations, client rules, and/or user-associated actions stored at one or more data stores **235**. Such data may influence communication routing to the devices within computing environment **214**, how data is stored or processed within computing environment **214**, among other actions.

(52) Notably, various other devices can further be used to influence communication routing and/or processing between devices within computing environment **214** and with devices outside of computing environment **214**. For example, as shown in FIG. 2, computing environment **214** may include a machine **240** that is a web server. Thus, computing environment **214** can retrieve data of interest, such as client information (e.g., product information, client rules, etc.), technical product details, news, current or predicted weather, and so on.

(53) In addition to computing environment **214** collecting data (e.g., as received from network devices, such as sensors, and client devices or other sources) to be processed as part of a big data analytics project, it may also receive data in real time as part of a streaming analytics environment. As noted, data may be collected using a variety of sources as communicated via different kinds of networks or locally. Such data may be received on a real-time streaming basis. For example, network devices may receive data periodically from network device sensors as the sensors continuously sense, monitor and track changes in their environments. Devices within computing environment **214** may also perform pre-analysis on data it receives to determine if the data received should be processed as part of an ongoing project. The data received and collected by computing environment **214**, no matter what the source or method or timing of receipt, may be processed over a period of time for a client to determine results data based on the client's needs and rules.

(54) FIG. 3 illustrates a representation of a conceptual model of a communications protocol system, according to embodiments of the present technology. More specifically, FIG. 3 identifies operation of a computing environment in an Open Systems Interaction model that corresponds to various connection components. The model **300** shows, for example, how a computing environment, such as computing environment **314** (or computing environment **214** in FIG. 2) may communicate with other devices in its network, and control how communications between the computing environment and other devices are executed and under what conditions.

(55) The model can include layers **301-307**. The layers are arranged in a stack. Each layer in the stack serves the layer one level higher than it (except for the application layer, which is the highest layer), and is served by the layer one level below it (except for the physical layer, which is the lowest layer). The physical layer is the lowest layer because it receives and transmits raw bites of data, and is the farthest layer from the user in a communications system. On the other hand, the application layer is the highest layer because it interacts directly with a software application.

(56) As noted, the model includes a physical layer **301**. Physical layer **301** represents physical communication, and can define parameters of that physical communication. For example, such physical communication may come in the form of electrical, optical, or electromagnetic signals. Physical layer **301** also defines protocols that may control communications within a data transmission network.

(57) Link layer **302** defines links and mechanisms used to transmit (i.e., move) data across a network. The link layer **302** manages node-to-node communications, such as within a grid computing environment. Link layer **302** can detect and correct errors (e.g., transmission errors in

the physical layer **301**) Link layer **302** can also include a media access control (MAC) layer and logical link control (LLC) layer.

(58) Network layer **303** defines the protocol for routing within a network. In other words, the network layer coordinates transferring data across nodes in a same network (e.g., such as a grid computing environment). Network layer **303** can also define the processes used to structure local addressing within the network.

(59) Transport layer **304** can manage the transmission of data and the quality of the transmission and/or receipt of that data. Transport layer **304** can provide a protocol for transferring data, such as, for example, a Transmission Control Protocol (TCP). Transport layer **304** can assemble and disassemble data frames for transmission. The transport layer can also detect transmission errors occurring in the layers below it.

(60) Session layer **305** can establish, maintain, and manage communication connections between devices on a network. In other words, the session layer controls the dialogues or nature of communications between network devices on the network. The session layer may also establish checkpointing, adjournment, termination, and restart procedures.

(61) Presentation layer **306** can provide translation for communications between the application and network layers. In other words, this layer may encrypt, decrypt and/or format data based on data types and/or encodings known to be accepted by an application or network layer.

(62) Application layer **307** interacts directly with software applications and end users, and manages communications between them. Application layer **307** can identify destinations, local resource states or availability and/or communication content or formatting using the applications.

(63) Intra-network connection components **321** and **322** are shown to operate in lower levels, such as physical layer **301** and link layer **302**, respectively. For example, a hub can operate in the physical layer, a switch can operate in the link layer, and a router can operate in the network layer. Inter-network connection components **323** and **328** are shown to operate on higher levels, such as layers **303-307**. For example, routers can operate in the network layer and network devices can operate in the transport, session, presentation, and application layers.

(64) As noted, a computing environment **314** can interact with and/or operate on, in various embodiments, one, more, all or any of the various layers. For example, computing environment **314** can interact with a hub (e.g., via the link layer) so as to adjust which devices the hub communicates with. The physical layer may be served by the link layer, so it may implement such data from the link layer. For example, the computing environment **314** may control which devices it will receive data from. For example, if the computing environment **314** knows that a certain network device has turned off, broken, or otherwise become unavailable or unreliable, the computing environment **314** may instruct the hub to prevent any data from being transmitted to the computing environment **314** from that network device. Such a process may be beneficial to avoid receiving data that is inaccurate or that has been influenced by an uncontrolled environment. As another example, computing environment **314** can communicate with a bridge, switch, router or gateway and influence which device within the system (e.g., system **200**) the component selects as a destination. In some embodiments, computing environment **314** can interact with various layers by exchanging communications with equipment operating on a particular layer by routing or modifying existing communications. In another embodiment, such as in a grid computing environment, a node may determine how data within the environment should be routed (e.g., which node should receive certain data) based on certain parameters or information provided by other layers within the model.

(65) As noted, the computing environment **314** may be a part of a communications grid environment, the communications of which may be implemented as shown in the protocol of FIG. 3. For example, referring back to FIG. 2, one or more of machines **220** and **240** may be part of a communications grid computing environment. A gridded computing environment may be employed in a distributed system with non-interactive workloads where data resides in memory on the machines, or compute nodes. In such an environment, analytic code, instead of a database

management system, controls the processing performed by the nodes. Data is co-located by pre-distributing it to the grid nodes, and the analytic code on each node loads the local data into memory. Each node may be assigned a particular task such as a portion of a processing project, or to organize or control other nodes within the grid.

(66) FIG. 4 illustrates a communications grid computing system **400** including a variety of control and worker nodes, according to embodiments of the present technology. Communications grid computing system **400** includes three control nodes and one or more worker nodes.

Communications grid computing system **400** includes control nodes **402**, **404**, and **406**. The control nodes are communicatively connected via communication paths **451**, **453**, and **455**. Therefore, the control nodes may transmit information (e.g., related to the communications grid or notifications), to and receive information from each other. Although communications grid computing system **400** is shown in FIG. 4 as including three control nodes, the communications grid may include more or less than three control nodes.

(67) Communications grid computing system (or just “communications grid”) **400** also includes one or more worker nodes. Shown in FIG. 4 are six worker nodes **410-420**. Although FIG. 4 shows six worker nodes, a communications grid according to embodiments of the present technology may include more or less than six worker nodes. The number of worker nodes included in a communications grid may be dependent upon how large the project or data set is being processed by the communications grid, the capacity of each worker node, the time designated for the communications grid to complete the project, among others. Each worker node within the communications grid **400** may be connected (wired or wirelessly, and directly or indirectly) to control nodes **402-406**. Therefore, each worker node may receive information from the control nodes (e.g., an instruction to perform work on a project) and may transmit information to the control nodes (e.g., a result from work performed on a project). Furthermore, worker nodes may communicate with each other (either directly or indirectly). For example, worker nodes may transmit data between each other related to a job being performed or an individual task within a job being performed by that worker node. However, in certain embodiments, worker nodes may not, for example, be connected (communicatively or otherwise) to certain other worker nodes. In an embodiment, worker nodes may only be able to communicate with the control node that controls it, and may not be able to communicate with other worker nodes in the communications grid, whether they are other worker nodes controlled by the control node that controls the worker node, or worker nodes that are controlled by other control nodes in the communications grid.

(68) A control node may connect with an external device with which the control node may communicate (e.g., a grid user, such as a server or computer, may connect to a controller of the grid). For example, a server or computer may connect to control nodes and may transmit a project or job to the node. The project may include a data set. The data set may be of any size. Once the control node receives such a project including a large data set, the control node may distribute the data set or projects related to the data set to be performed by worker nodes. Alternatively, for a project including a large data set, the data set may be received or stored by a machine other than a control node (e.g., a HADOOP® standard-compliant data node employing the HADOOP® Distributed File System, or HDFS).

(69) Control nodes may maintain knowledge of the status of the nodes in the grid (i.e., grid status information), accept work requests from clients, subdivide the work across worker nodes, and coordinate the worker nodes, among other responsibilities. Worker nodes may accept work requests from a control node and provide the control node with results of the work performed by the worker node. A grid may be started from a single node (e.g., a machine, computer, server, etc.). This first node may be assigned or may start as the primary control node that will control any additional nodes that enter the grid.

(70) When a project is submitted for execution (e.g., by a client or a controller of the grid) it may be assigned to a set of nodes. After the nodes are assigned to a project, a data structure (i.e., a

communicator) may be created. The communicator may be used by the project for information to be shared between the project codes running on each node. A communication handle may be created on each node. A handle, for example, is a reference to the communicator that is valid within a single process on a single node, and the handle may be used when requesting communications between nodes.

(71) A control node, such as control node **402**, may be designated as the primary control node. A server, computer or other external device may connect to the primary control node. Once the control node receives a project, the primary control node may distribute portions of the project to its worker nodes for execution. For example, when a project is initiated on communications grid **400**, primary control node **402** controls the work to be performed for the project in order to complete the project as requested or instructed. The primary control node may distribute work to the worker nodes based on various factors, such as which subsets or portions of projects may be completed most efficiently and in the correct amount of time. For example, a worker node may perform analysis on a portion of data that is already local to (e.g., stored on) the worker node. The primary control node also coordinates and processes the results of the work performed by each worker node after each worker node executes and completes its job. For example, the primary control node may receive a result from one or more worker nodes, and the control node may organize (e.g., collect and assemble) the results received and compile them to produce a complete result for the project received from the end user.

(72) Any remaining control nodes, such as control nodes **404** and **406**, may be assigned as backup control nodes for the project. In an embodiment, backup control nodes may not control any portion of the project. Instead, backup control nodes may serve as a backup for the primary control node and take over as primary control node if the primary control node were to fail. If a communications grid were to include only a single control node, and the control node were to fail (e.g., the control node is shut off or breaks), then the communications grid as a whole may fail and any project or job being run on the communications grid may fail and may not complete. While the project may be run again, such a failure may cause a delay (severe delay in some cases, such as overnight delay) in completion of the project. Therefore, a grid with multiple control nodes, including a backup control node, may be beneficial.

(73) To add another node or machine to the grid, the primary control node may open a pair of listening sockets, for example. A socket may be used to accept work requests from clients, and the second socket may be used to accept connections from other grid nodes. The primary control node may be provided with a list of other nodes (e.g., other machines, computers, servers) that will participate in the grid, and the role that each node will fill in the grid. Upon startup of the primary control node (e.g., the first node on the grid), the primary control node may use a network protocol to start the server process on every other node in the grid. Command line parameters, for example, may inform each node of one or more pieces of information, such as: the role that the node will have in the grid, the host name of the primary control node, and the port number on which the primary control node is accepting connections from peer nodes, among others. The information may also be provided in a configuration file, transmitted over a secure shell tunnel, or received from a configuration server, among others. While the other machines in the grid may not initially know about the configuration of the grid, that information may also be sent to each other node by the primary control node. Updates of the grid information may also be subsequently sent to those nodes.

(74) For any control node other than the primary control node added to the grid, the control node may open three sockets. The first socket may accept work requests from clients, the second socket may accept connections from other grid members, and the third socket may connect (e.g., permanently) to the primary control node. When a control node (e.g., primary control node) receives a connection from another control node, it first checks to see if the peer node is in the list of configured nodes in the grid. If it is not on the list, the control node may clear the connection. If

it is on the list, it may then attempt to authenticate the connection. If authentication is successful, the authenticating node may transmit information to its peer, such as the port number on which a node is listening for connections, the host name of the node, and information about how to authenticate the node, among other information. When a node, such as the new control node, receives information about another active node, it will check to see if it already has a connection to that other node. If it does not have a connection to that node, it may then establish a connection to that control node.

(75) Any worker node added to the grid may establish a connection to the primary control node and any other control nodes on the grid. After establishing the connection, it may authenticate itself to the grid (e.g., any control nodes, including both primary and backup, or a server or user controlling the grid). After successful authentication, the worker node may accept configuration information from the control node.

(76) When a node joins a communications grid (e.g., when the node is powered on or connected to an existing node on the grid or both), the node is assigned (e.g., by an operating system of the grid) a universally unique identifier (UUID). This unique identifier may help other nodes and external entities (devices, users, etc.) to identify the node and distinguish it from other nodes. When a node is connected to the grid, the node may share its unique identifier with the other nodes in the grid. Since each node may share its unique identifier, each node may know the unique identifier of every other node on the grid. Unique identifiers may also designate a hierarchy of each of the nodes (e.g., backup control nodes) within the grid. For example, the unique identifiers of each of the backup control nodes may be stored in a list of backup control nodes to indicate an order in which the backup control nodes will take over for a failed primary control node to become a new primary control node. However, a hierarchy of nodes may also be determined using methods other than using the unique identifiers of the nodes. For example, the hierarchy may be predetermined, or may be assigned based on other predetermined factors.

(77) The grid may add new machines at any time (e.g., initiated from any control node). Upon adding a new node to the grid, the control node may first add the new node to its table of grid nodes. The control node may also then notify every other control node about the new node. The nodes receiving the notification may acknowledge that they have updated their configuration information.

(78) Primary control node **402** may, for example, transmit one or more communications to backup control nodes **404** and **406** (and, for example, to other control or worker nodes within the communications grid). Such communications may be sent periodically, at fixed time intervals, between known fixed stages of the project's execution, among other protocols. The communications transmitted by primary control node **402** may be of varied types and may include a variety of types of information. For example, primary control node **402** may transmit snapshots (e.g., status information) of the communications grid so that backup control node **404** always has a recent snapshot of the communications grid. The snapshot or grid status may include, for example, the structure of the grid (including, for example, the worker nodes in the grid, unique identifiers of the nodes, or their relationships with the primary control node) and the status of a project (including, for example, the status of each worker node's portion of the project). The snapshot may also include analysis or results received from worker nodes in the communications grid. The backup control nodes may receive and store the backup data received from the primary control node. The backup control nodes may transmit a request for such a snapshot (or other information) from the primary control node, or the primary control node may send such information periodically to the backup control nodes.

(79) As noted, the backup data may allow the backup control node to take over as primary control node if the primary control node fails without requiring the grid to start the project over from scratch. If the primary control node fails, the backup control node that will take over as primary control node may retrieve the most recent version of the snapshot received from the primary

control node and use the snapshot to continue the project from the stage of the project indicated by the backup data. This may prevent failure of the project as a whole.

(80) A backup control node may use various methods to determine that the primary control node has failed. In one example of such a method, the primary control node may transmit (e.g., periodically) a communication to the backup control node that indicates that the primary control node is working and has not failed, such as a heartbeat communication. The backup control node may determine that the primary control node has failed if the backup control node has not received a heartbeat communication for a certain predetermined period of time. Alternatively, a backup control node may also receive a communication from the primary control node itself (before it failed) or from a worker node that the primary control node has failed, for example because the primary control node has failed to communicate with the worker node.

(81) Different methods may be performed to determine which backup control node of a set of backup control nodes (e.g., backup control nodes **404** and **406**) will take over for failed primary control node **402** and become the new primary control node. For example, the new primary control node may be chosen based on a ranking or “hierarchy” of backup control nodes based on their unique identifiers. In an alternative embodiment, a backup control node may be assigned to be the new primary control node by another device in the communications grid or from an external device (e.g., a system infrastructure or an end user, such as a server or computer, controlling the communications grid). In another alternative embodiment, the backup control node that takes over as the new primary control node may be designated based on bandwidth or other statistics about the communications grid.

(82) A worker node within the communications grid may also fail. If a worker node fails, work being performed by the failed worker node may be redistributed amongst the operational worker nodes. In an alternative embodiment, the primary control node may transmit a communication to each of the operable worker nodes still on the communications grid that each of the worker nodes should purposefully fail also. After each of the worker nodes fail, they may each retrieve their most recent saved checkpoint of their status and re-start the project from that checkpoint to minimize lost progress on the project being executed.

(83) FIG. 5 illustrates a flow chart showing an example process **500** for adjusting a communications grid or a work project in a communications grid after a failure of a node, according to embodiments of the present technology. The process may include, for example, receiving grid status information including a project status of a portion of a project being executed by a node in the communications grid, as described in operation **502**. For example, a control node (e.g., a backup control node connected to a primary control node and a worker node on a communications grid) may receive grid status information, where the grid status information includes a project status of the primary control node or a project status of the worker node. The project status of the primary control node and the project status of the worker node may include a status of one or more portions of a project being executed by the primary and worker nodes in the communications grid. The process may also include storing the grid status information, as described in operation **504**. For example, a control node (e.g., a backup control node) may store the received grid status information locally within the control node. Alternatively, the grid status information may be sent to another device for storage where the control node may have access to the information.

(84) The process may also include receiving a failure communication corresponding to a node in the communications grid in operation **506**. For example, a node may receive a failure communication including an indication that the primary control node has failed, prompting a backup control node to take over for the primary control node. In an alternative embodiment, a node may receive a failure that a worker node has failed, prompting a control node to reassign the work being performed by the worker node. The process may also include reassigning a node or a portion of the project being executed by the failed node, as described in operation **508**. For

example, a control node may designate the backup control node as a new primary control node based on the failure communication upon receiving the failure communication. If the failed node is a worker node, a control node may identify a project status of the failed worker node using the snapshot of the communications grid, where the project status of the failed worker node includes a status of a portion of the project being executed by the failed worker node at the failure time.

(85) The process may also include receiving updated grid status information based on the reassignment, as described in operation **510**, and transmitting a set of instructions based on the updated grid status information to one or more nodes in the communications grid, as described in operation **512**. The updated grid status information may include an updated project status of the primary control node or an updated project status of the worker node. The updated information may be transmitted to the other nodes in the grid to update their stale stored information.

(86) FIG. **6** illustrates a portion of a communications grid computing system **600** including a control node and a worker node, according to embodiments of the present technology. Communications grid computing system **600** includes one control node (control node **602**) and one worker node (worker node **610**) for purposes of illustration, but may include more worker and/or control nodes. The control node **602** is communicatively connected to worker node **610** via communication path **650**. Therefore, control node **602** may transmit information (e.g., related to the communications grid or notifications), to and receive information from worker node **610** via path **650**.

(87) Similar to in FIG. **4**, communications grid computing system (or just “communications grid”) **600** includes data processing nodes (control node **602** and worker node **610**). Nodes **602** and **610** include multi-core data processors. Each node **602** and **610** includes a grid-enabled software component (GESC) **620** that executes on the data processor associated with that node and interfaces with buffer memory **622** also associated with that node. Each node **602** and **610** includes database management software (DBMS) **628** that executes on a database server (not shown) at control node **602** and on a database server (not shown) at worker node **610**.

(88) Each node also includes a data store **624**. Data stores **624**, similar to network-attached data stores **110** in FIG. **1** and data stores **235** in FIG. **2**, are used to store data to be processed by the nodes in the computing environment. Data stores **624** may also store any intermediate or final data generated by the computing system after being processed, for example in non-volatile memory. However in certain embodiments, the configuration of the grid computing environment allows its operations to be performed such that intermediate and final data results can be stored solely in volatile memory (e.g., RAM), without a requirement that intermediate or final data results be stored to non-volatile types of memory. Storing such data in volatile memory may be useful in certain situations, such as when the grid receives queries (e.g., ad hoc) from a client and when responses, which are generated by processing large amounts of data, need to be generated quickly or on-the-fly. In such a situation, the grid may be configured to retain the data within memory so that responses can be generated at different levels of detail and so that a client may interactively query against this information.

(89) Each node also includes a user-defined function (UDF) **626**. The UDF provides a mechanism for the DBMS **628** to transfer data to or receive data from the database stored in the data stores **624** that are managed by the DBMS **628**. For example, UDF **626** can be invoked by the DBMS **628** to provide data to the GESC **620** for processing. The UDF **626** may establish a socket connection (not shown) with the GESC **620** to transfer the data. Alternatively, the UDF **626** can transfer data to the GESC **620** by writing data to shared memory accessible by both the UDF **626** and the GESC **620**.

(90) The GESC **620** at the nodes **602** and **610** may be connected via a network, such as network **108** shown in FIG. **1**. Therefore, nodes **602** and **610** can communicate with each other via the network using a predetermined communication protocol such as, for example, the Message Passing Interface (MPI). Each GESC **620** can engage in point-to-point communication with the GESC at another node or in collective communication with multiple GESCs via the network. The GESC **620**

at each node may contain identical (or nearly identical) software instructions. Each node may be capable of operating as either a control node or a worker node. The GESC at the control node **602** can communicate, over a communication path **652**, with a client device **630**. More specifically, control node **602** may communicate with client application **632** hosted by the client device **630** to receive queries and to respond to those queries after processing large amounts of data.

(91) DBMS **628** may control the creation, maintenance, and use of database or data structure (not shown) within a nodes **602** or **610**. The database may organize data stored in data stores **624**. The DBMS **628** at control node **602** may accept requests for data and transfer the appropriate data for the request. With such a process, collections of data may be distributed across multiple physical locations. In this example, each node **602** and **610** stores a portion of the total data managed by the management system in its associated data store **624**.

(92) Furthermore, the DBMS may be responsible for protecting against data loss using replication techniques. Replication includes providing a backup copy of data stored on one node on one or more other nodes. Therefore, if one node fails, the data from the failed node can be recovered from a replicated copy residing at another node. However, as described herein with respect to FIG. **4**, data or status information for each node in the communications grid may also be shared with each node on the grid.

(93) FIG. **7** illustrates a flow chart showing an example method **700** for executing a project within a grid computing system, according to embodiments of the present technology. As described with respect to FIG. **6**, the GESC at the control node may transmit data with a client device (e.g., client device **630**) to receive queries for executing a project and to respond to those queries after large amounts of data have been processed. The query may be transmitted to the control node, where the query may include a request for executing a project, as described in operation **702**. The query can contain instructions on the type of data analysis to be performed in the project and whether the project should be executed using the grid-based computing environment, as shown in operation **704**.

(94) To initiate the project, the control node may determine if the query requests use of the grid-based computing environment to execute the project. If the determination is no, then the control node initiates execution of the project in a solo environment (e.g., at the control node), as described in operation **710**. If the determination is yes, the control node may initiate execution of the project in the grid-based computing environment, as described in operation **706**. In such a situation, the request may include a requested configuration of the grid. For example, the request may include a number of control nodes and a number of worker nodes to be used in the grid when executing the project. After the project has been completed, the control node may transmit results of the analysis yielded by the grid, as described in operation **708**. Whether the project is executed in a solo or grid-based environment, the control node provides the results of the project, as described in operation **712**.

(95) As noted with respect to FIG. **2**, the computing environments described herein may collect data (e.g., as received from network devices, such as sensors, such as network devices **204-209** in FIG. **2**, and client devices or other sources) to be processed as part of a data analytics project, and data may be received in real time as part of a streaming analytics environment (e.g., ESP). Data may be collected using a variety of sources as communicated via different kinds of networks or locally, such as on a real-time streaming basis. For example, network devices may receive data periodically from network device sensors as the sensors continuously sense, monitor and track changes in their environments. More specifically, an increasing number of distributed applications develop or produce continuously flowing data from distributed sources by applying queries to the data before distributing the data to geographically distributed recipients. An event stream processing engine (ESPE) may continuously apply the queries to the data as it is received and determines which entities should receive the data. Client or other devices may also subscribe to the ESPE or other devices processing ESP data so that they can receive data after processing, based on

for example the entities determined by the processing engine. For example, client devices **230** in FIG. **2** may subscribe to the ESPE in computing environment **214**. In another example, event subscription devices **1024a-c**, described further with respect to FIG. **10**, may also subscribe to the ESPE. The ESPE may determine or define how input data or event streams from network devices or other publishers (e.g., network devices **204-209** in FIG. **2**) are transformed into meaningful output data to be consumed by subscribers, such as for example client devices **230** in FIG. **2**.

(96) FIG. **8** illustrates a block diagram including components of an Event Stream Processing Engine (ESPE), according to embodiments of the present technology. ESPE **800** may include one or more projects **802**. A project may be described as a second-level container in an engine model managed by ESPE **800** where a thread pool size for the project may be defined by a user. Each project of the one or more projects **802** may include one or more continuous queries **804** that contain data flows, which are data transformations of incoming event streams. The one or more continuous queries **804** may include one or more source windows **806** and one or more derived windows **808**.

(97) The ESPE may receive streaming data over a period of time related to certain events, such as events or other data sensed by one or more network devices. The ESPE may perform operations associated with processing data created by the one or more devices. For example, the ESPE may receive data from the one or more network devices **204-209** shown in FIG. **2**. As noted, the network devices may include sensors that sense different aspects of their environments, and may collect data over time based on those sensed observations. For example, the ESPE may be implemented within one or more of machines **220** and **240** shown in FIG. **2**. The ESPE may be implemented within such a machine by an ESP application. An ESP application may embed an ESPE with its own dedicated thread pool or pools into its application space where the main application thread can do application-specific work and the ESPE processes event streams at least by creating an instance of a model into processing objects.

(98) The engine container is the top-level container in a model that manages the resources of the one or more projects **802**. In an illustrative embodiment, for example, there may be only one ESPE **800** for each instance of the ESP application, and ESPE **800** may have a unique engine name. Additionally, the one or more projects **802** may each have unique project names, and each query may have a unique continuous query name and begin with a uniquely named source window of the one or more source windows **806**. ESPE **800** may or may not be persistent.

(99) Continuous query modeling involves defining directed graphs of windows for event stream manipulation and transformation. A window in the context of event stream manipulation and transformation is a processing node in an event stream processing model. A window in a continuous query can perform aggregations, computations, pattern-matching, and other operations on data flowing through the window. A continuous query may be described as a directed graph of source, relational, pattern matching, and procedural windows. The one or more source windows **806** and the one or more derived windows **808** represent continuously executing queries that generate updates to a query result set as new event blocks stream through ESPE **800**. A directed graph, for example, is a set of nodes connected by edges, where the edges have a direction associated with them.

(100) An event object may be described as a packet of data accessible as a collection of fields, with at least one of the fields defined as a key or unique identifier (ID). The event object may be created using a variety of formats including binary, alphanumeric, XML, etc. Each event object may include one or more fields designated as a primary identifier (ID) for the event so ESPE **800** can support operation codes (opcodes) for events including insert, update, upsert, and delete. Upsert opcodes update the event if the key field already exists; otherwise, the event is inserted. For illustration, an event object may be a packed binary representation of a set of field values and include both metadata and field data associated with an event. The metadata may include an opcode indicating if the event represents an insert, update, delete, or upsert, a set of flags indicating if the

event is a normal, partial-update, or a retention generated event from retention policy management, and a set of microsecond timestamps that can be used for latency measurements.

(101) An event block object may be described as a grouping or package of event objects. An event stream may be described as a flow of event block objects. A continuous query of the one or more continuous queries **804** transforms a source event stream made up of streaming event block objects published into ESPE **800** into one or more output event streams using the one or more source windows **806** and the one or more derived windows **808**. A continuous query can also be thought of as data flow modeling.

(102) The one or more source windows **806** are at the top of the directed graph and have no windows feeding into them. Event streams are published into the one or more source windows **806**, and from there, the event streams may be directed to the next set of connected windows as defined by the directed graph. The one or more derived windows **808** are all instantiated windows that are not source windows and that have other windows streaming events into them. The one or more derived windows **808** may perform computations or transformations on the incoming event streams. The one or more derived windows **808** transform event streams based on the window type (that is operators such as join, filter, compute, aggregate, copy, pattern match, procedural, union, etc.) and window settings. As event streams are published into ESPE **800**, they are continuously queried, and the resulting sets of derived windows in these queries are continuously updated.

(103) FIG. **9** illustrates a flow chart showing an example process including operations performed by an event stream processing engine, according to some embodiments of the present technology. As noted, the ESPE **800** (or an associated ESP application) defines how input event streams are transformed into meaningful output event streams. More specifically, the ESP application may define how input event streams from publishers (e.g., network devices providing sensed data) are transformed into meaningful output event streams consumed by subscribers (e.g., a data analytics project being executed by a machine or set of machines).

(104) Within the application, a user may interact with one or more user interface windows presented to the user in a display under control of the ESPE independently or through a browser application in an order selectable by the user. For example, a user may execute an ESP application, which causes presentation of a first user interface window, which may include a plurality of menus and selectors such as drop down menus, buttons, text boxes, hyperlinks, etc. associated with the ESP application as understood by a person of skill in the art. As further understood by a person of skill in the art, various operations may be performed in parallel, for example, using a plurality of threads.

(105) At operation **900**, an ESP application may define and start an ESPE, thereby instantiating an ESPE at a device, such as machine **220** and/or **240**. In an operation **902**, the engine container is created. For illustration, ESPE **800** may be instantiated using a function call that specifies the engine container as a manager for the model.

(106) In an operation **904**, the one or more continuous queries **804** are instantiated by ESPE **800** as a model. The one or more continuous queries **804** may be instantiated with a dedicated thread pool or pools that generate updates as new events stream through ESPE **800**. For illustration, the one or more continuous queries **804** may be created to model business processing logic within ESPE **800**, to predict events within ESPE **800**, to model a physical system within ESPE **800**, to predict the physical system state within ESPE **800**, etc. For example, as noted, ESPE **800** may be used to support sensor data monitoring and management (e.g., sensing may include force, torque, load, strain, position, temperature, air pressure, fluid flow, chemical properties, resistance, electromagnetic fields, radiation, irradiance, proximity, acoustics, moisture, distance, speed, vibrations, acceleration, electrical potential, or electrical current, etc.).

(107) ESPE **800** may analyze and process events in motion or “event streams.” Instead of storing data and running queries against the stored data, ESPE **800** may store queries and stream data through them to allow continuous analysis of data as it is received. The one or more source

windows **806** and the one or more derived windows **808** may be created based on the relational, pattern matching, and procedural algorithms that transform the input event streams into the output event streams to model, simulate, score, test, predict, etc. based on the continuous query model defined and application to the streamed data.

(108) In an operation **906**, a publish/subscribe (pub/sub) capability is initialized for ESPE **800**. In an illustrative embodiment, a pub/sub capability is initialized for each project of the one or more projects **802**. To initialize and enable pub/sub capability for ESPE **800**, a port number may be provided. Pub/sub clients can use a host name of an ESP device running the ESPE and the port number to establish pub/sub connections to ESPE **800**.

(109) FIG. **10** illustrates an ESP system **1000** interfacing between publishing device **1022** and event subscribing devices **1024a-c**, according to embodiments of the present technology. ESP system **1000** may include ESP device or subsystem **1001**, event publishing device **1022**, an event subscribing device A **1024a**, an event subscribing device B **1024b**, and an event subscribing device C **1024c**. Input event streams are output to ESP subsystem **1001** by publishing device **1022**. In alternative embodiments, the input event streams may be created by a plurality of publishing devices. The plurality of publishing devices further may publish event streams to other ESP devices. The one or more continuous queries instantiated by ESPE **800** may analyze and process the input event streams to form output event streams output to event subscribing device A **1024a**, event subscribing device B **1024b**, and event subscribing device C **1024c**. ESP system **1000** may include a greater or a fewer number of event subscribing devices of event subscribing devices.

(110) Publish-subscribe is a message-oriented interaction paradigm based on indirect addressing. Processed data recipients specify their interest in receiving information from ESPE **800** by subscribing to specific classes of events, while information sources publish events to ESPE **800** without directly addressing the receiving parties. ESPE **800** coordinates the interactions and processes the data. In some cases, the data source receives confirmation that the published information has been received by a data recipient.

(111) A publish/subscribe API may be described as a library that enables an event publisher, such as publishing device **1022**, to publish event streams into ESPE **800** or an event subscriber, such as event subscribing device A **1024a**, event subscribing device B **1024b**, and event subscribing device C **1024c**, to subscribe to event streams from ESPE **800**. For illustration, one or more publish/subscribe APIs may be defined. Using the publish/subscribe API, an event publishing application may publish event streams into a running event stream processor project source window of ESPE **800**, and the event subscription application may subscribe to an event stream processor project source window of ESPE **800**.

(112) The publish/subscribe API provides cross-platform connectivity and endianness compatibility between ESP application and other networked applications, such as event publishing applications instantiated at publishing device **1022**, and event subscription applications instantiated at one or more of event subscribing device A **1024a**, event subscribing device B **1024b**, and event subscribing device C **1024c**.

(113) Referring back to FIG. **9**, operation **906** initializes the publish/subscribe capability of ESPE **800**. In an operation **908**, the one or more projects **802** are started. The one or more started projects may run in the background on an ESP device. In an operation **910**, an event block object is received from one or more computing device of the event publishing device **1022**.

(114) ESP subsystem **1001** may include a publishing client **1002**, ESPE **800**, a subscribing client A **1004**, a subscribing client B **1006**, and a subscribing client C **1008**. Publishing client **1002** may be started by an event publishing application executing at publishing device **1022** using the publish/subscribe API. Subscribing client A **1004** may be started by an event subscription application A, executing at event subscribing device A **1024a** using the publish/subscribe API. Subscribing client B **1006** may be started by an event subscription application B executing at event subscribing device B **1024b** using the publish/subscribe API. Subscribing client C **1008** may be

started by an event subscription application C executing at event subscribing device C **1024c** using the publish/subscribe API.

(115) An event block object containing one or more event objects is injected into a source window of the one or more source windows **806** from an instance of an event publishing application on event publishing device **1022**. The event block object may be generated, for example, by the event publishing application and may be received by publishing client **1002**. A unique ID may be maintained as the event block object is passed between the one or more source windows **806** and/or the one or more derived windows **808** of ESPE **800**, and to subscribing client A **1004**, subscribing client B **1006**, and subscribing client C **1008** and to event subscription device A **1024a**, event subscription device B **1024b**, and event subscription device C **1024c**. Publishing client **1002** may further generate and include a unique embedded transaction ID in the event block object as the event block object is processed by a continuous query, as well as the unique ID that publishing device **1022** assigned to the event block object.

(116) In an operation **912**, the event block object is processed through the one or more continuous queries **804**. In an operation **914**, the processed event block object is output to one or more computing devices of the event subscribing devices **1024a-c**. For example, subscribing client A **1004**, subscribing client B **1006**, and subscribing client C **1008** may send the received event block object to event subscription device A **1024a**, event subscription device B **1024b**, and event subscription device C **1024c**, respectively.

(117) ESPE **800** maintains the event block containership aspect of the received event blocks from when the event block is published into a source window and works its way through the directed graph defined by the one or more continuous queries **804** with the various event translations before being output to subscribers. Subscribers can correlate a group of subscribed events back to a group of published events by comparing the unique ID of the event block object that a publisher, such as publishing device **1022**, attached to the event block object with the event block ID received by the subscriber.

(118) In an operation **916**, a determination is made concerning whether or not processing is stopped. If processing is not stopped, processing continues in operation **910** to continue receiving the one or more event streams containing event block objects from the, for example, one or more network devices. If processing is stopped, processing continues in an operation **918**. In operation **918**, the started projects are stopped. In operation **920**, the ESPE is shutdown.

(119) As noted, in some embodiments, big data is processed for an analytics project after the data is received and stored. In other embodiments, distributed applications process continuously flowing data in real-time from distributed sources by applying queries to the data before distributing the data to geographically distributed recipients. As noted, an event stream processing engine (ESPE) may continuously apply the queries to the data as it is received and determines which entities receive the processed data. This allows for large amounts of data being received and/or collected in a variety of environments to be processed and distributed in real time. For example, as shown with respect to FIG. 2, data may be collected from network devices that may include devices within the internet of things, such as devices within a home automation network. However, such data may be collected from a variety of different resources in a variety of different environments. In any such situation, embodiments of the present technology allow for real-time processing of such data.

(120) Aspects of the current disclosure provide technical solutions to technical problems, such as computing problems that arise when an ESP device fails which results in a complete service interruption and potentially significant data loss. The data loss can be catastrophic when the streamed data is supporting mission critical operations such as those in support of an ongoing manufacturing or drilling operation. An embodiment of an ESP system achieves a rapid and seamless failover of ESPE running at the plurality of ESP devices without service interruption or data loss, thus significantly improving the reliability of an operational system that relies on the live or real-time processing of the data streams. The event publishing systems, the event subscribing

systems, and each ESPE not executing at a failed ESP device are not aware of or effected by the failed ESP device. The ESP system may include thousands of event publishing systems and event subscribing systems. The ESP system keeps the failover logic and awareness within the boundaries of out-messaging network connector and out-messaging network device.

(121) In one example embodiment, a system is provided to support a failover when event stream processing (ESP) event blocks. The system includes, but is not limited to, an out-messaging network device and a computing device. The computing device includes, but is not limited to, a processor and a computer-readable medium operably coupled to the processor. The processor is configured to execute an ESP engine (ESPE). The computer-readable medium has instructions stored thereon that, when executed by the processor, cause the computing device to support the failover. An event block object is received from the ESPE that includes a unique identifier. A first status of the computing device as active or standby is determined. When the first status is active, a second status of the computing device as newly active or not newly active is determined. Newly active is determined when the computing device is switched from a standby status to an active status. When the second status is newly active, a last published event block object identifier that uniquely identifies a last published event block object is determined. A next event block object is selected from a non-transitory computer-readable medium accessible by the computing device. The next event block object has an event block object identifier that is greater than the determined last published event block object identifier. The selected next event block object is published to an out-messaging network device. When the second status of the computing device is not newly active, the received event block object is published to the out-messaging network device. When the first status of the computing device is standby, the received event block object is stored in the non-transitory computer-readable medium.

(122) FIG. 11 is a flow chart of an example of a process for generating and using a machine-learning model according to some aspects. Machine learning is a branch of artificial intelligence that relates to mathematical models that can learn from, categorize, and make predictions about data. Such mathematical models, which can be referred to as machine-learning models, can classify input data among two or more classes; cluster input data among two or more groups; predict a result based on input data; identify patterns or trends in input data; identify a distribution of input data in a space; or any combination of these. Examples of machine-learning models can include (i) neural networks; (ii) decision trees, such as classification trees and regression trees; (iii) classifiers, such as Naïve bias classifiers, logistic regression classifiers, ridge regression classifiers, random forest classifiers, least absolute shrinkage and selector (LASSO) classifiers, and support vector machines; (iv) clusterers, such as k-means clusterers, mean-shift clusterers, and spectral clusterers; (v) factorizers, such as factorization machines, principal component analyzers and kernel principal component analyzers; and (vi) ensembles or other combinations of machine-learning models. In some examples, neural networks can include deep neural networks, feed-forward neural networks, recurrent neural networks, convolutional neural networks, radial basis function (RBF) neural networks, echo state neural networks, long short-term memory neural networks, bi-directional recurrent neural networks, gated neural networks, hierarchical recurrent neural networks, stochastic neural networks, modular neural networks, spiking neural networks, dynamic neural networks, cascading neural networks, neuro-fuzzy neural networks, or any combination of these.

(123) Different machine-learning models may be used interchangeably to perform a task. Examples of tasks that can be performed at least partially using machine-learning models include various types of scoring; bioinformatics; cheminformatics; software engineering; fraud detection; customer segmentation; generating online recommendations; adaptive websites; determining customer lifetime value; search engines; placing advertisements in real time or near real time; classifying DNA sequences; affective computing; performing natural language processing and understanding; object recognition and computer vision; robotic locomotion; playing games; optimization and metaheuristics; detecting network intrusions; medical diagnosis and monitoring; or predicting when

an asset, such as a machine, will need maintenance.

(124) Any number and combination of tools can be used to create machine-learning models. Examples of tools for creating and managing machine-learning models can include SAS® Enterprise Miner, SAS® Rapid Predictive Modeler, and SAS® Model Manager, SAS Cloud Analytic Services (CAS)®, SAS Viya® of all which are by SAS Institute Inc. of Cary, North Carolina.

(125) Machine-learning models can be constructed through an at least partially automated (e.g., with little or no human involvement) process called training. During training, input data can be iteratively supplied to a machine-learning model to enable the machine-learning model to identify patterns related to the input data or to identify relationships between the input data and output data. With training, the machine-learning model can be transformed from an untrained state to a trained state. Input data can be split into one or more training sets and one or more validation sets, and the training process may be repeated multiple times. The splitting may follow a k-fold cross-validation rule, a leave-one-out-rule, a leave-p-out rule, or a holdout rule. An overview of training and using a machine-learning model is described below with respect to the flow chart of FIG. 11.

(126) In block **1102**, training data is received. In some examples, the training data is received from a remote database or a local database, constructed from various subsets of data, or input by a user. The training data can be used in its raw form for training a machine-learning model or pre-processed into another form, which can then be used for training the machine-learning model. For example, the raw form of the training data can be smoothed, truncated, aggregated, clustered, or otherwise manipulated into another form, which can then be used for training the machine-learning model.

(127) In block **1104**, a machine-learning model is trained using the training data. The machine-learning model can be trained in a supervised, unsupervised, or semi-supervised manner. In supervised training, each input in the training data is correlated to a desired output. This desired output may be a scalar, a vector, or a different type of data structure such as text or an image. This may enable the machine-learning model to learn a mapping between the inputs and desired outputs. In unsupervised training, the training data includes inputs, but not desired outputs, so that the machine-learning model has to find structure in the inputs on its own. In semi-supervised training, only some of the inputs in the training data are correlated to desired outputs.

(128) In block **1106**, the machine-learning model is evaluated. For example, an evaluation dataset can be obtained, for example, via user input or from a database. The evaluation dataset can include inputs correlated to desired outputs. The inputs can be provided to the machine-learning model and the outputs from the machine-learning model can be compared to the desired outputs. If the outputs from the machine-learning model closely correspond with the desired outputs, the machine-learning model may have a high degree of accuracy. For example, if 90% or more of the outputs from the machine-learning model are the same as the desired outputs in the evaluation dataset, the machine-learning model may have a high degree of accuracy. Otherwise, the machine-learning model may have a low degree of accuracy. The 90% number is an example only. A realistic and desirable accuracy percentage is dependent on the problem and the data.

(129) In some examples, if, at block **1108**, the machine-learning model has an inadequate degree of accuracy for a particular task, the process can return to block **1104**, where the machine-learning model can be further trained using additional training data or otherwise modified to improve accuracy. However, if, at block **1108**, the machine-learning model has an adequate degree of accuracy for the particular task, the process can continue to block **1110**.

(130) In block **1110**, new data is received. In some examples, the new data is received from a remote database or a local database, constructed from various subsets of data, or input by a user. The new data may be unknown to the machine-learning model. For example, the machine-learning model may not have previously processed or analyzed the new data.

(131) In block **1112**, the trained machine-learning model is used to analyze the new data and

provide a result. For example, the new data can be provided as input to the trained machine-learning model. The trained machine-learning model can analyze the new data and provide a result that includes a classification of the new data into a particular class, a clustering of the new data into a particular group, a prediction based on the new data, or any combination of these.

(132) In block **1114**, the result is post-processed. For example, the result can be added to, multiplied with, or otherwise combined with other data as part of a job. As another example, the result can be transformed from a first format, such as a time series format, into another format, such as a count series format. Any number and combination of operations can be performed on the result during post-processing.

(133) A more specific example of a machine-learning model is the neural network **1200** shown in FIG. **12**. The neural network **1200** is represented as multiple layers of neurons **1208** that can exchange data between one another via connections **1255** that may be selectively instantiated thereamong. The layers include an input layer **1202** for receiving input data provided at inputs **1222**, one or more hidden layers **1204**, and an output layer **1206** for providing a result at outputs **1277**. The hidden layer(s) **1204** are referred to as hidden because they may not be directly observable or have their inputs or outputs directly accessible during the normal functioning of the neural network **1200**. Although the neural network **1200** is shown as having a specific number of layers and neurons for exemplary purposes, the neural network **1200** can have any number and combination of layers, and each layer can have any number and combination of neurons.

(134) The neurons **1208** and connections **1255** thereamong may have numeric weights, which can be tuned during training of the neural network **1200**. For example, training data can be provided to at least the inputs **1222** to the input layer **1202** of the neural network **1200**, and the neural network **1200** can use the training data to tune one or more numeric weights of the neural network **1200**. In some examples, the neural network **1200** can be trained using backpropagation. Backpropagation can include determining a gradient of a particular numeric weight based on a difference between an actual output of the neural network **1200** at the outputs **1277** and a desired output of the neural network **1200**. Based on the gradient, one or more numeric weights of the neural network **1200** can be updated to reduce the difference therebetween, thereby increasing the accuracy of the neural network **1200**. This process can be repeated multiple times to train the neural network **1200**. For example, this process can be repeated hundreds or thousands of times to train the neural network **1200**.

(135) In some examples, the neural network **1200** is a feed-forward neural network. In a feed-forward neural network, the connections **1255** are instantiated and/or weighted so that every neuron **1208** only propagates an output value to a subsequent layer of the neural network **1200**. For example, data may only move one direction (forward) from one neuron **1208** to the next neuron **1208** in a feed-forward neural network. Such a “forward” direction may be defined as proceeding from the input layer **1202** through the one or more hidden layers **1204**, and toward the output layer **1206**.

(136) In other examples, the neural network **1200** may be a recurrent neural network. A recurrent neural network can include one or more feedback loops among the connections **1255**, thereby allowing data to propagate in both forward and backward through the neural network **1200**. Such a “backward” direction may be defined as proceeding in the opposite direction of forward, such as from the output layer **1206** through the one or more hidden layers **1204**, and toward the input layer **1202**. This can allow for information to persist within the recurrent neural network. For example, a recurrent neural network can determine an output based at least partially on information that the recurrent neural network has seen before, giving the recurrent neural network the ability to use previous input to inform the output.

(137) In some examples, the neural network **1200** operates by receiving a vector of numbers from one layer; transforming the vector of numbers into a new vector of numbers using a matrix of numeric weights, a nonlinearity, or both; and providing the new vector of numbers to a subsequent

layer (“subsequent” in the sense of moving “forward”) of the neural network **1200**. Each subsequent layer of the neural network **1200** can repeat this process until the neural network **1200** outputs a final result at the outputs **1277** of the output layer **1206**. For example, the neural network **1200** can receive a vector of numbers at the inputs **1222** of the input layer **1202**. The neural network **1200** can multiply the vector of numbers by a matrix of numeric weights to determine a weighted vector. The matrix of numeric weights can be tuned during the training of the neural network **1200**. The neural network **1200** can transform the weighted vector using a nonlinearity, such as a sigmoid tangent or the hyperbolic tangent. In some examples, the nonlinearity can include a rectified linear unit, which can be expressed using the equation $y=\max(x, 0)$ where y is the output and x is an input value from the weighted vector. The transformed output can be supplied to a subsequent layer (e.g., a hidden layer **1204**) of the neural network **1200**. The subsequent layer of the neural network **1200** can receive the transformed output, multiply the transformed output by a matrix of numeric weights and a nonlinearity, and provide the result to yet another layer of the neural network **1200** (e.g., another, subsequent, hidden layer **1204**). This process continues until the neural network **1200** outputs a final result at the outputs **1277** of the output layer **1206**.

(138) As also depicted in FIG. **12**, the neural network **1200** may be implemented either through the execution of the instructions of one or more routines **1244** by central processing units (CPUs), or through the use of one or more neuromorphic devices **1250** that incorporate a set of memristors (or other similar components) that each function to implement one of the neurons **1208** in hardware. Where multiple neuromorphic devices **1250** are used, they may be interconnected in a depth-wise manner to enable implementing neural networks with greater quantities of layers, and/or in a width-wise manner to enable implementing neural networks having greater quantities of neurons **1208** per layer.

(139) The neuromorphic device **1250** may incorporate a storage interface **1299** by which neural network configuration data **1293** that is descriptive of various parameters and hyperparameters of the neural network **1200** may be stored and/or retrieved. More specifically, the neural network configuration data **1293** may include such parameters as weighting and/or biasing values derived through the training of the neural network **1200**, as has been described. Alternatively or additionally, the neural network configuration data **1293** may include such hyperparameters as the manner in which the neurons **1208** are to be interconnected (e.g., feed-forward or recurrent), the trigger function to be implemented within the neurons **1208**, the quantity of layers and/or the overall quantity of the neurons **1208**. The neural network configuration data **1293** may provide such information for more than one neuromorphic device **1250** where multiple ones have been interconnected to support larger neural networks.

(140) Other examples of the present disclosure may include any number and combination of machine-learning models having any number and combination of characteristics. The machine-learning model(s) can be trained in a supervised, semi-supervised, or unsupervised manner, or any combination of these. The machine-learning model(s) can be implemented using a single computing device or multiple computing devices, such as the communications grid computing system **400** discussed above.

(141) Implementing some examples of the present disclosure at least in part by using machine-learning models can reduce the total number of processing iterations, time, memory, electrical power, or any combination of these consumed by a computing device when analyzing data. For example, a neural network may more readily identify patterns in data than other approaches. This may enable the neural network to analyze the data using fewer processing cycles and less memory than other approaches, while obtaining a similar or greater level of accuracy.

(142) Some machine-learning approaches may be more efficiently and speedily executed and processed with machine-learning specific processors (e.g., not a generic CPU). Such processors may also provide an energy savings when compared to generic CPUs. For example, some of these processors can include a graphical processing unit (GPU), an application-specific integrated circuit

(ASIC), a field-programmable gate array (FPGA), an artificial intelligence (AI) accelerator, a neural computing core, a neural computing engine, a neural processing unit, a purpose-built chip architecture for deep learning, and/or some other machine-learning specific processor that implements a machine learning approach or one or more neural networks using semiconductor (e.g., silicon (Si), gallium arsenide(GaAs)) devices. These processors may also be employed in heterogeneous computing architectures with a number of and/or a variety of different types of cores, engines, nodes, and/or layers to achieve various energy efficiencies, processing speed improvements, data communication speed improvements, and/or data efficiency targets and improvements throughout various parts of the system when compared to a homogeneous computing architecture that employs CPUs for general purpose computing.

(143) FIG. 13 illustrates various aspects of the use of containers 1336 as a mechanism to allocate processing, storage and/or other resources of a processing system 1300 to the performance of various analyses. More specifically, in a processing system 1300 that includes one or more node devices 1330 (e.g., the aforementioned grid system 400), the processing, storage and/or other resources of each node device 1330 may be allocated through the instantiation and/or maintenance of multiple containers 1336 within the node devices 1330 to support the performance(s) of one or more analyses. As each container 1336 is instantiated, predetermined amounts of processing, storage and/or other resources may be allocated thereto as part of creating an execution environment therein in which one or more executable routines 1334 may be executed to cause the performance of part or all of each analysis that is requested to be performed.

(144) It may be that at least a subset of the containers 1336 are each allocated a similar combination and amounts of resources so that each is of a similar configuration with a similar range of capabilities, and therefore, are interchangeable. This may be done in embodiments in which it is desired to have at least such a subset of the containers 1336 already instantiated prior to the receipt of requests to perform analyses, and thus, prior to the specific resource requirements of each of those analyses being known.

(145) Alternatively or additionally, it may be that at least a subset of the containers 1336 are not instantiated until after the processing system 1300 receives requests to perform analyses where each request may include indications of the resources required for one of those analyses. Such information concerning resource requirements may then be used to guide the selection of resources and/or the amount of each resource allocated to each such container 1336. As a result, it may be that one or more of the containers 1336 are caused to have somewhat specialized configurations such that there may be differing types of containers to support the performance of different analyses and/or different portions of analyses.

(146) It may be that the entirety of the logic of a requested analysis is implemented within a single executable routine 1334. In such embodiments, it may be that the entirety of that analysis is performed within a single container 1336 as that single executable routine 1334 is executed therein. However, it may be that such a single executable routine 1334, when executed, is at least intended to cause the instantiation of multiple instances of itself that are intended to be executed at least partially in parallel. This may result in the execution of multiple instances of such an executable routine 1334 within a single container 1336 and/or across multiple containers 1336.

(147) Alternatively or additionally, it may be that the logic of a requested analysis is implemented with multiple differing executable routines 1334. In such embodiments, it may be that at least a subset of such differing executable routines 1334 are executed within a single container 1336. However, it may be that the execution of at least a subset of such differing executable routines 1334 is distributed across multiple containers 1336.

(148) Where an executable routine 1334 of an analysis is under development, and/or is under scrutiny to confirm its functionality, it may be that the container 1336 within which that executable routine 1334 is to be executed is additionally configured assist in limiting and/or monitoring aspects of the functionality of that executable routine 1334. More specifically, the execution

environment provided by such a container **1336** may be configured to enforce limitations on accesses that are allowed to be made to memory and/or I/O addresses to control what storage locations and/or I/O devices may be accessible to that executable routine **1334**. Such limitations may be derived based on comments within the programming code of the executable routine **1334** and/or other information that describes what functionality the executable routine **1334** is expected to have, including what memory and/or I/O accesses are expected to be made when the executable routine **1334** is executed. Then, when the executable routine **1334** is executed within such a container **1336**, the accesses that are attempted to be made by the executable routine **1334** may be monitored to identify any behavior that deviates from what is expected.

(149) Where the possibility exists that different executable routines **1334** may be written in different programming languages, it may be that different subsets of containers **1336** are configured to support different programming languages. In such embodiments, it may be that each executable routine **1334** is analyzed to identify what programming language it is written in, and then what container **1336** is assigned to support the execution of that executable routine **1334** may be at least partially based on the identified programming language. Where the possibility exists that a single requested analysis may be based on the execution of multiple executable routines **1334** that may each be written in a different programming language, it may be that at least a subset of the containers **1336** are configured to support the performance of various data structure and/or data format conversion operations to enable a data object output by one executable routine **1334** written in one programming language to be accepted as an input to another executable routine **1334** written in another programming language.

(150) As depicted, at least a subset of the containers **1336** may be instantiated within one or more VMs **1331** that may be instantiated within one or more node devices **1330**. Thus, in some embodiments, it may be that the processing, storage and/or other resources of at least one node device **1330** may be partially allocated through the instantiation of one or more VMs **1331**, and then in turn, may be further allocated within at least one VM **1331** through the instantiation of one or more containers **1336**.

(151) In some embodiments, it may be that such a nested allocation of resources may be carried out to effect an allocation of resources based on two differing criteria. By way of example, it may be that the instantiation of VMs **1331** is used to allocate the resources of a node device **1330** to multiple users or groups of users in accordance with any of a variety of service agreements by which amounts of processing, storage and/or other resources are paid for each such user or group of users. Then, within each VM **1331** or set of VMs **1331** that is allocated to a particular user or group of users, containers **1336** may be allocated to distribute the resources allocated to each VM **1331** among various analyses that are requested to be performed by that particular user or group of users.

(152) As depicted, where the processing system **1300** includes more than one node device **1330**, the processing system **1300** may also include at least one control device **1350** within which one or more control routines **1354** may be executed to control various aspects of the use of the node device(s) **1330** to perform requested analyses. By way of example, it may be that at least one control routine **1354** implements logic to control the allocation of the processing, storage and/or other resources of each node device **1330** to each VM **1331** and/or container **1336** that is instantiated therein. Thus, it may be the control device(s) **1350** that effects a nested allocation of resources, such as the aforementioned example allocation of resources based on two differing criteria.

(153) As also depicted, the processing system **1300** may also include one or more distinct requesting devices **1370** from which requests to perform analyses may be received by the control device(s) **1350**. Thus, and by way of example, it may be that at least one control routine **1354** implements logic to monitor for the receipt of requests from authorized users and/or groups of users for various analyses to be performed using the processing, storage and/or other resources of the node device(s) **1330** of the processing system **1300**. The control device(s) **1350** may receive

indications of the availability of resources, the status of the performances of analyses that are already underway, and/or still other status information from the node device(s) **1330** in response to polling, at a recurring interval of time, and/or in response to the occurrence of various preselected events. More specifically, the control device(s) **1350** may receive indications of status for each container **1336**, each VM **1331** and/or each node device **1330**. At least one control routine **1354** may implement logic that may use such information to select container(s) **1336**, VM(s) **1331** and/or node device(s) **1330** that are to be used in the execution of the executable routine(s) **1334** associated with each requested analysis.

(154) As further depicted, in some embodiments, the one or more control routines **1354** may be executed within one or more containers **1356** and/or within one or more VMs **1351** that may be instantiated within the one or more control devices **1350**. It may be that multiple instances of one or more varieties of control routine **1354** may be executed within separate containers **1356**, within separate VMs **1351** and/or within separate control devices **1350** to better enable parallelized control over parallel performances of requested analyses, to provide improved redundancy against failures for such control functions, and/or to separate differing ones of the control routines **1354** that perform different functions. By way of example, it may be that multiple instances of a first variety of control routine **1354** that communicate with the requesting device(s) **1370** are executed in a first set of containers **1356** instantiated within a first VM **1351**, while multiple instances of a second variety of control routine **1354** that control the allocation of resources of the node device(s) **1330** are executed in a second set of containers **1356** instantiated within a second VM **1351**. It may be that the control of the allocation of resources for performing requested analyses may include deriving an order of performance of portions of each requested analysis based on such factors as data dependencies thereamong, as well as allocating the use of containers **1336** in a manner that effectuates such a derived order of performance.

(155) Where multiple instances of control routine **1354** are used to control the allocation of resources for performing requested analyses, such as the assignment of individual ones of the containers **1336** to be used in executing executable routines **1334** of each of multiple requested analyses, it may be that each requested analysis is assigned to be controlled by just one of the instances of control routine **1354**. This may be done as part of treating each requested analysis as one or more “ACID transactions” that each have the four properties of atomicity, consistency, isolation and durability such that a single instance of control routine **1354** is given full control over the entirety of each such transaction to better ensure that all of each such transaction is either entirely performed or is entirely not performed. Allowing partial performances to occur may cause cache incoherencies and/or data corruption issues.

(156) As additionally depicted, the control device(s) **1350** may communicate with the requesting device(s) **1370** and with the node device(s) **1330** through portions of a network **1399** extending thereamong. Again, such a network as the depicted network **1399** may be based on any of a variety of wired and/or wireless technologies, and may employ any of a variety of protocols by which commands, status, data and/or still other varieties of information may be exchanged. It may be that one or more instances of a control routine **1354** cause the instantiation and maintenance of a web portal or other variety of portal that is based on any of a variety of communication protocols, etc. (e.g., a restful API). Through such a portal, requests for the performance of various analyses may be received from requesting device(s) **1370**, and/or the results of such requested analyses may be provided thereto. Alternatively or additionally, it may be that one or more instances of a control routine **1354** cause the instantiation of and maintenance of a message passing interface and/or message queues. Through such an interface and/or queues, individual containers **1336** may each be assigned to execute at least one executable routine **1334** associated with a requested analysis to cause the performance of at least a portion of that analysis.

(157) Although not specifically depicted, it may be that at least one control routine **1354** may include logic to implement a form of management of the containers **1336** based on the Kubernetes

container management platform promulgated by Cloud Native Computing Foundation of San Francisco, CA, USA. In such embodiments, containers **1336** in which executable routines **1334** of requested analyses may be instantiated within “pods” (not specifically shown) in which other containers may also be instantiated for the execution of other supporting routines. Such supporting routines may cooperate with control routine(s) **1354** to implement a communications protocol with the control device(s) **1350** via the network **1399** (e.g., a message passing interface, one or more message queues, etc.). Alternatively or additionally, such supporting routines may serve to provide access to one or more storage repositories (not specifically shown) in which at least data objects may be stored for use in performing the requested analyses.

(158) FIG. **14** shows a block diagram of an example of a system **1400** for automatically transforming data into an electronic data structure according to some aspects of the present disclosure. The system **1400** includes a client device **1404** associated with a user **1406**. Examples of the client device **1404** can include a laptop computer, desktop computer, mobile telephone, tablet, e-reader, smartwatch, etc. The user **1406** may operate the client device **1404** to draft a structured document **1422** in a human-readable declarative language, such as XML or JSON. For instance, the user **1406** may use an editor of a graphical user interface **1432** displayed on the client device **1404** to draft the structured document **1422**. In the structured document **1422**, the user **1406** can include one or more conditions **1424** usable to customize how a dataset **1412** is to be transformed into one or more electronic data structures **1426**. A portion of the dataset **1412** that satisfies the one or more conditions **1424** may be included in the one or more electronic data structures **1426**, while a remainder of the dataset **1412** may be excluded from the one or more electronic data structures **1426**.

(159) There can be any number of conditions **1424** included in the structured document **1422**. One example of such a condition may be that only data entries in the dataset **1412** in which a particular variable has a particular value (e.g., “X=3”) should be included in the electronic data structures **1426**. Another example of such a condition may be that only data entries in the dataset **1412** in which a particular type of variable exists should be included in the electronic data structures **1426**. Still another example of such a condition can be a date range. Data entries with timestamps inside the date range may be included in an electronic data structure **1426**, while data entries with timestamps outside the date range may be excluded from the electronic data structure **1426**. Alternatively, data entries with timestamps inside the date range may be assigned to a first electronic data structure, while data entries with timestamps outside the date range may be assigned to a second electronic data structure. The conditions **1424** may be used alone or in combination with each other to control which subset of the dataset **1412** to include in the one or more electronic data structures **1426**. In some cases, only data entries that satisfy two or more of (e.g., all of) the conditions **1424** may be included in the electronic data structure(s) **1426**. In other cases, any data entries that satisfy at least one of the conditions **1424** may be included in the electronic data structure(s) **1426**.

(160) After drafting the structured document **1422**, the user **1406** can transmit the structured document **1422** to a server system **1402**, which may include one or more servers, desktop computers, etc. For example, the user **1406** can operate the client device **1404** to submit the structured document **1422** to the server system **1402** via one or more networks **1408**, such as a LAN or the Internet. The user **1406** may also transmit an indicator of the dataset **1412** to be processed using the structured document **1422**. In some examples, the indicator of the dataset **1412** may be a name of the dataset **1412** or other unique identifier of the dataset **1412**.

(161) The server system **1402** can receive the structured document **1422** and the indicator of the dataset **1412**, for example, as part of a request from the client device **1404**. In response to receiving the request, the server system **1402** can retrieve dataset **1412** based on the indicator. For example, the server system **1402** can transmit a query **1414** that includes the indicator to the database **1410**. In response to the query **1414**, the database **1410** can return the dataset **1412** to the server system

1402.

(162) The server system **1402** can next apply an interpreter program **1420** to the dataset **1412** to process the dataset **1412**. In some examples, the server system **1402** may have access to multiple interpreter programs **1418** (e.g., stored in a repository **1416**), where each interpreter program may be configured to process a different type of input dataset and/or generate a different type of output electronic data structure. The server system **1402** can select a particular interpreter program **1420** from among the multiple available interpreter programs **1418** for use in processing the dataset **1412**. The server system **1402** may make this selection based on the structured document **1422**, the dataset **1412**, or both. For example, the server system **1402** may make this selection based on the current format of the input dataset **1412** and the target format for the output electronic data structure(s) **1426**. The target format for the output electronic data structure(s) **1426** may be defined in the structured document **1422** by the user **1406**. For instance, the server system **1402** can analyze the dataset **1412** and the structured document **1422** to determine a current format of the input dataset **1412** and a target format for an output electronic data structure **1426**, respectively, and then select a particular interpreter program **1420** that is compatible with those two formats. The input/output formats associated with each interpreter program may be stored in a lookup table, which can be accessed by the server system **1402** to facilitate making this selection.

(163) After selecting the interpreter program **1420**, the server system **1402** can retrieve (e.g., download from the repository **1416**) and execute the interpreter program **1420**. When executed, the interpreter program **1420** can ingest the dataset **1412** and the structured document **1422**. The interpreter program **1420** can apply the one or more conditions **1424** defined in the structured document **1422** to the dataset **1412**, to thereby produce one or more electronic data structures **1426** that satisfy those conditions. In particular, the interpreter program **1420** can identify a subset of the dataset **1412** that satisfies the one or more conditions **1424** and incorporate that subset into the one or more electronic data structures **1426**. The interpreter program **1420** may exclude a remainder of the dataset **1412** from the one or more electronic data structures **1426**. In generating the one or more electronic data structures **1426**, the interpreter program **1420** can configure the one or more electronic data structures **1426** to have the target format defined in the structured document **1422**. The target format of the electronic data structure(s) **1426** can be different from the current format of the dataset **1412**, such that a format transformation takes place in generating the electronic data structure(s) **1426**. For instance, the dataset **1412** can be in a single table, whereas the electronic data structures **1426** can include two or more tables across which the dataset **1412** is divided based on the conditions **1424**. In some examples, the electronic data structure(s) **1426** can be electronic spreadsheets that can be exported as one or more files **1430** for viewing or other usage by the user **1406**. Those files may then be transmitted back to the client device **1404**, where they may be used for any number of purposes. In some examples, the content of the files **1430** may be displayed in the GUI **1432** to the user.

(164) Depending on the size of the dataset **1412**, the electronic data structures **1426** may still be very large (e.g., gigabytes or terabytes in size). In some cases, such as with sparse data, a significant amount of the fields in the electronic data structures **1426** may be empty and, thus, unnecessarily consuming storage space. To reduce this unnecessary storage consumption, the server system **1402** may execute a compaction process on the one or more electronic data structures **1426** to thereby produce one or more compacted data structures **1428**. The compaction process may involve removing (e.g., deleting) empty fields from the one or more electronic data structures **1426**. For instance, the compaction process may involve removing empty columns and empty rows from the one or more electronic data structures **1426**. An empty column can be a column for which all of the fields are empty, and an empty row can be a row for which all of the fields are empty. Eliminating the empty fields from the one or more electronic data structures **1426** can free up the memory regions associated with those empty fields for other purposes, allowing for a more efficient use of memory space. This may be critically important in certain scenarios where the

electronic data structures **1426** will be used with devices containing a very limited amount of memory space, such as on edge devices. In some example, the compacted data structure(s) **1428** can be exported as one or more files **1430** for viewing or other usage by the user **1406**. Those files may then be transmitted back to the client device **1404**, where they may be used for any number of purposes. In some examples, the content of the files **1430** may be displayed in the GUI **1432** to the user.

(165) Turning now to FIG. **15**, shown is a flowchart of an example of a process for automatically transforming data into an electronic data structure according to some aspects of the present disclosure. Other examples may include more operations, fewer operations, different operations, or a different order of operations than is shown in FIG. **15**. The operations of FIG. **15** are described below with reference to the components of FIG. **14** described above.

(166) In block **1502**, a computing system (e.g., server system **1402**) receives a structured document **1422** that describes one or more conditions **1424** selected to customize one or more electronic data structures **1426**. The one or more conditions **1424** are defined in the structured document **1422** using a declarative language. The computing system can receive the structured document **1422** from the user (e.g., by the user writing the structured document in an editor program) or from another source, such as a repository of stored structured documents.

(167) In block **1504**, the computing system queries a database **1410** to retrieve data (e.g., dataset **1412**) from the database **1410**. For example, the computing system can transmit a query **1414** to the database **1410**. The query **1414** can include a unique identifier of the data to be retrieved. In response to receiving the query **1414**, the database **1410** can retrieve the data associated with the unique identifier and provide it to the computing system.

(168) In block **1506**, the computing system selects an interpreter program **1420**, from among a plurality of interpreter programs **1418**, based on the structured document **1422** and the data. For example, the computing system can determine a first format of the data and a second format for the electronic data structures **1426**. The computing system can determine the first format by analyzing the data, and can determine the second format by analyzing the structured document **1422**. The computing system can then select whichever of the interpreter programs is configured to operate on data in the first format and produce electronic data structures **1426** in the second format. The computing system can determine which of the interpreter programs are configured to operate on data in the first format and produce electronic data structures **1426** in the second format by referencing a lookup table, which may map either of the interpreter programs to its input and output formats.

(169) In block **1508**, the computing system executes the interpreter program **1420** to identify a subset of the data that satisfies the one or more conditions **1424**. For example, the computing system can apply each of the conditions **1424** to each data entry in the data to determine which portions (e.g., data entries) of the data satisfy the one or more conditions **1424**.

(170) In block **1510**, the computing system (e.g., the interpreter program **1420** executing thereon) generates the one or more electronic data structures **1426** based on the subset of the data. For instance, the computing system generates the at least two electronic data structures **1426** based on the subset of the data.

(171) The one or more electronic data structures **1426** can include the subset and exclude a remainder of the data. The one or more electronic data structures **1426** may be configured as in-memory data structures stored in random access memory or other volatile memory. Alternatively, the one or more electronic data structures **1426** may be configured as one or more files, such as spreadsheet files or other types of files. In some such examples, a single file may include multiple electronic data structures. For instance, a single spreadsheet file may include multiple separate spreadsheets that correspond to multiple electronic data structures.

(172) Because the format of the data may be different from the format of the one or more electronic data structures **1426**, there may be a format conversion that takes place in block **510** to produce the

one or more electronic data structures **1426**. The logic for implementing this conversion may be pre-programmed into the particular interpreter program **1420** by a developer thereof.

(173) In some examples, the electronic data structures **1426** can be different from one another in format, content, and/or structure. For instance, based on the subset of the data, the interpreter program **1420** can generate three electronic data structures that differ format, content, and/or structure from one another. Thus, a single interpreter program may be capable of dividing a single dataset into multiple electronic data structures with different characteristics from one another.

(174) In block **1512**, the computing system compacts the one or more electronic data structures **1426** to generate one or more compacted data structures **1428**. This compaction process may involve removing empty fields from the one or more electronic data structures **1426**, thereby reducing their size in memory.

(175) In block **1514**, the computing system can provide the one or more compacted data structures **1428** to a user **1406**. For example, the computing system can output the one or more compacted data structures **1428** to the user **1406** via a graphical user interface **1432**. In some examples, the graphical user interface **1432** may be output on a display of the computing system. Alternatively, the graphical user interface **1432** may be output on a display of a client device **1404** interacting with the computing system. The compacted data structures **1428** may be output as one or more tables in the graphical user interface **1432**.

(176) Turning now to FIG. **16**, shown is an example of data **1600** that can be transformed into a custom electronic data structure according to some aspects of the present disclosure. In this example, the data **1600** is in the form of a table, though other formats are possible. The table includes N columns and M rows. The first column **1602** can correspond to different variables, while a remainder of the columns can correspond to different types of values associated with those variables. As shown, many of the data fields (e.g., cells) in the table can be empty. After the data **1600** is converted into a target type of electronic data structure, the electronic data structure may undergo a compaction process to remove some or all of these empty data fields. This may help to conserve valuable memory space.

(177) FIG. **17** shows an example of a structured document **1700** usable to transform data into a custom electronic data structure according to some aspects of the present disclosure. In this example, the structured document **1700** includes a first segment **1702** that defines a first electronic data structure (e.g., a first spreadsheet) to be output by the interpreter program. The first electronic data structure can be titled “IRs” and only uses the variable names as conditions. The structures document **1700** also includes a second segment **1704** that defines a second electronic data structure (e.g., a second spreadsheet). The second electronic data structure can be titled “Curves & Groups” and uses other forms of conditions shown in subsections **1706**, **1708**. Subsections **1706**, **1708** are alternative condition options that can be selected. Subsection **1706** allows for any data entries in the columns titled “Type” and “Role” to be included in the second electronic data structure if their values are equal to “Curve” and “ir,” respectively. Subsection **1708** allows for any data entries in the column titled “Type” to be included in the second electronic data structure if their values equal “Group.” It will be appreciated that both of these examples are intended to be illustrative and non-limiting. Other examples may involve other kinds of mathematical operators other than “equal to,” such as greater than, less than, greater than or equal to, less than or equal to, etc.

(178) Another example of a structured document **1800** is shown in FIG. **18**. This structured document **1800** shows an example of a condition **1802** involving a date range. In particular, the condition **1802** specifies that any data entries with timestamps between the start date of Dec. 2, 2023 and the end date of Dec. 3, 2023, and that correspond to a certain type of variable (“IR1”), should be included in the electronic data structure. Other types of conditions are also possible.

(179) Turning now to FIG. **19**, shown is an example of transforming data into multiple custom electronic data structures **1906a-n** according to some aspects of the present disclosure. In this example, the computing system executes an interpreter program **1420**, which can ingest a

structured document **1422** and a dataset **1412**. Based on the conditions defined in the structured document **1422**, the interpreter program **1420** can parse the dataset **1412** into any number of electronic data structures, such as the N electronic data structures **1906a-n** shown in FIG. **19**. Each of the electronic data structures **1906a-n** may have different data than the others, such that they do not having any overlapping data. Alternatively, two or more of the electronic data structures **1906a-n** may have at least some overlapping data to one another. Once generated, one or more of the electronic data structures **1906a-n** may undergo a compaction process before being transmitted to a user. An example of such a compaction process is shown in FIG. **20**.

(180) Referring now to FIG. **20**, shown is an example an electronic data structure **2000** that is a matrix containing M columns by N rows. As noted above, the electronic data structure **2000** may undergo a compaction process, in which rows and/or columns are removed. For example, rows and/or columns that only have empty fields (all of their fields are empty) may be removed. This may result in a compacted data structure **2002** that has X columns and Y rows, where X may be less than M, and where Y may be less than N. In the particular example shown in FIG. **20**, row **4** is an empty row because all of its cells are empty. Similarly, column **5** is an empty column because all of its cells are empty. Thus, during the compaction process, row **4** and column **5** may be removed from the electronic data structure **2000** to produce a compacted data structure **2002** of reduced size. Removing these excess rows from the electronic data structure **2000** can conserve memory space.

(181) In the previous description, for the purposes of explanation, specific details are set forth in order to provide a thorough understanding of examples of the technology. But various examples can be practiced without these specific details. The figures and description are not intended to be restrictive.

(182) The previous description provides examples that are not intended to limit the scope, applicability, or configuration of the disclosure. Rather, the previous description of the examples provides those skilled in the art with an enabling description for implementing an example. Various changes may be made in the function and arrangement of elements without departing from the spirit and scope of the technology as set forth in the appended claims.

(183) Specific details are given in the previous description to provide a thorough understanding of the examples. But the examples may be practiced without these specific details. For example, circuits, systems, networks, processes, and other components can be shown as components in block diagram form to prevent obscuring the examples in unnecessary detail. In other examples, well-known circuits, processes, algorithms, structures, and techniques may be shown without unnecessary detail in order to avoid obscuring the examples.

(184) Also, individual examples may have been described as a process that is depicted as a flowchart, a flow diagram, a data flow diagram, a structure diagram, or a block diagram. Although a flowchart can describe the operations as a sequential process, many of the operations can be performed in parallel or concurrently. In addition, the order of the operations can be re-arranged. And a process can have more or fewer operations than are depicted in a figure. A process can correspond to a method, a function, a procedure, a subroutine, a subprogram, etc. When a process corresponds to a function, its termination can correspond to a return of the function to the calling function or the main function.

(185) Systems depicted in some of the figures can be provided in various configurations. In some examples, the systems can be configured as a distributed system where one or more components of the system are distributed across one or more networks in a cloud computing system.

Claims

1. A method comprising: receiving, by one or more processors, a predefined structured document drafted by a user in a human-readable declarative language, wherein the predefined structured document is a transferrable file, wherein the structured document describes one or more conditions

selected by the user to customize one or more electronic data structures, wherein the one or more conditions are defined in the structured document using the human-readable declarative language, and wherein the one or more conditions specify variable values selected by the user for inclusion in the one or more electronic data structures; querying, by the one or more processors, a database to retrieve data from the database, wherein the querying is performed independently of the one or more conditions expressed in the predefined structured document; ingesting, by an interpreter program executing on the one or more processors, the structured document and the data; filtering, by the interpreter program, the data retrieved from the database based on the one or more conditions specified in the structured document to identify a subset of the data that satisfies the one or more conditions; generating, by the interpreter program executing on the one or more processors, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data; after generating the one or more electronic data structures, compacting, by the one or more processors, the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures; and providing, by the one or more processors, the one or more compacted data structures to the user.

2. The method of claim 1, further comprising: determining a first format of the data; determining, based on the structured document, a second format in which the data is to be organized in the one or more electronic data structures; and selecting the interpreter program from among a plurality of interpreter programs based on the first format of the data and the second format for the one or more electronic data structures, wherein each interpreter program of the plurality of interpreter programs is configured to process input data in a respective input format and output corresponding data structures in a respective output format.

3. The method of claim 1, wherein the one or more conditions include a particular value for a particular variable being present in the data.

4. The method of claim 1, wherein the one or more conditions include a particular type of variable being present in the data.

5. The method of claim 1, wherein the one or more conditions include a particular data range used for splitting the data into at least two electronic data structures.

6. The method of claim 1, wherein providing the one or more compacted data structures to the user involves exporting the one or more compacted data structures to a file as one or more spreadsheets.

7. The method of claim 1, wherein providing the one or more compacted data structures to the user involves outputting the one or more compacted data structures as one or more tables in a graphical user interface that is presented to the user.

8. The method of claim 1, wherein the one or more electronic data structures include at least two electronic spreadsheets.

9. The method of claim 1, wherein the human-readable declarative language is eXtensible Markup Language (XML) or JavaScript Object Notation (JSON).

10. The method of claim 1, wherein the data is organized in the one or more electronic data structures in a second format that is different from a first format of the data.

11. A non-transitory computer-readable medium comprising program code that is executable by one or more processors for causing the one or more processors to perform operations including: receiving a predefined structured document drafted by a user in a human-readable declarative language, wherein the predefined structured document is a transferrable file, wherein the structured document describes one or more conditions selected by the user to customize one or more electronic data structures, wherein the one or more conditions are defined in the structured document using the human-readable declarative language, and wherein the one or more conditions specify variable values selected by the user for inclusion in the one or more electronic data structures; querying a database to retrieve data from the database, wherein the querying is

performed independently of the one or more conditions expressed in the predefined structured document; ingesting, by an interpreter program, the structured document and the data; filtering, by the interpreter program, the data retrieved from the database based on the one or more conditions specified in the structured document to identify a subset of the data that satisfies the one or more conditions; generating, by the interpreter program, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data; after generating the one or more electronic data structures, compacting the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures; and providing the one or more compacted data structures to the user.

12. The non-transitory computer-readable medium of claim 11, wherein the operations further comprise: determining a first format of the data; determining, based on the structured document, a second format in which the data is to be organized in the one or more electronic data structures; and selecting the interpreter program from among a plurality of interpreter programs based on the first format of the data and the second format for the one or more electronic data structures, wherein each interpreter program of the plurality of interpreter programs is configured to process input data in a respective input format and output corresponding data structures in a respective output format.

13. The non-transitory computer-readable medium of claim 11, wherein the one or more conditions include a particular value for a particular variable being present in the data.

14. The non-transitory computer-readable medium of claim 11, wherein the one or more conditions include a particular type of variable being present in the data.

15. The non-transitory computer-readable medium of claim 11, wherein the one or more conditions include a particular data range used for splitting the data into at least two electronic data structures.

16. The non-transitory computer-readable medium of claim 11, wherein providing the one or more compacted data structures to the user involves exporting the one or more compacted data structures to a file as one or more spreadsheets.

17. The non-transitory computer-readable medium of claim 11, wherein providing the one or more compacted data structures to the user involves outputting the one or more compacted data structures as one or more tables in a graphical user interface that is presented to the user.

18. The non-transitory computer-readable medium of claim 11, wherein the one or more electronic data structures include at least two electronic spreadsheets, and wherein the empty fields include empty rows or empty columns of the at least two electronic spreadsheets.

19. The non-transitory computer-readable medium of claim 11, wherein the human-readable declarative language is eXtensible Markup Language (XML) or JavaScript Object Notation (JSON).

20. The non-transitory computer-readable medium of claim 11, wherein the data is organized in the one or more electronic data structures in a second format that is different from a first format of the data.

21. A system comprising: one or more processors; and one or more memories comprising program code that is executable by the one or more processors for causing the one or more processors to perform operations including: receiving a predefined structured document drafted by a user in a human-readable declarative language, wherein the predefined structured document is a transferrable file, wherein the structured document describes one or more conditions selected by the user to customize one or more electronic data structures, wherein the one or more conditions are defined in the structured document using the human-readable declarative language, and wherein the one or more conditions specify variable values selected by the user for inclusion in the one or more electronic data structures; querying a database to retrieve data from the database, wherein the querying is performed independently of the one or more conditions expressed in the predefined structured document; ingesting, by an interpreter program, the structured document and the data; filtering, by the interpreter program, the data retrieved from the database based on the one or more

conditions specified in the structured document to identify a subset of the data that satisfies the one or more conditions; generating, by the interpreter program, the one or more electronic data structures based on the subset of the data, wherein the one or more electronic data structures include the identified subset of the data and exclude a remainder of the data; after generating the one or more electronic data structures, compacting the one or more electronic data structures by removing empty fields from the one or more electronic data structures, to thereby generate one or more compacted data structures; and providing the one or more compacted data structures to the user.

22. The system of claim 21, wherein the operations further comprise: determining a first format of the data; determining, based on the structured document, a second format in which the data is to be organized in the one or more electronic data structures; and selecting the interpreter program from among a plurality of interpreter programs based on the first format of the data and the second format for the one or more electronic data structures, wherein each interpreter program of the plurality of interpreter programs is configured to process input data in a respective input format and output corresponding data structures in a respective output format.

23. The system of claim 21, wherein the one or more conditions include a particular value for a particular variable being present in the data.

24. The system of claim 21, wherein the one or more conditions include a particular type of variable being present in the data.

25. The system of claim 21, wherein the one or more conditions include a particular data range used for splitting the data into at least two electronic data structures.

26. The system of claim 21, wherein providing the one or more compacted data structures to the user involves exporting the one or more compacted data structures to a file as one or more spreadsheets.

27. The system of claim 21, wherein providing the one or more compacted data structures to the user involves outputting the one or more compacted data structures as one or more tables in a graphical user interface that is presented to the user.

28. The system of claim 21, wherein the one or more electronic data structures include at least two electronic spreadsheets.

29. The system of claim 21, wherein the human-readable declarative language is eXtensible Markup Language (XML) or JavaScript Object Notation (JSON).

30. The system of claim 21, wherein the data is organized in the one or more electronic data structures in a second format that is different from a first format of the data.
