

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12405972
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Portisch; Jan et al.

Systems and methods for manipulating time-dependent relationships in knowledge graph data structures

Abstract

Various examples are directed to systems and methods for utilizing relationship data. A computing system may receive a time-dependent query against a knowledge graph data structure. The computing system may access confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true. The computing system may determine that at least one of a beginning or an end of a first time period associated with a test relationship of the time-dependent query is not defined by the knowledge graph data structure. The computing system may determine a response to the first time-dependent query indicating a veracity of the test relationship at a test point-in-time using the first plurality of confirmation points-and-time.

Inventors:	Portisch; Jan (Bruchsal, DE), Bracholdt; Sandra (Dielheim, DE)
Applicant:	SAP SE (Walldorf, DE)
Family ID:	91281984
Assignee:	SAP SE (Walldorf, DE)
Appl. No.:	17/973311
Filed:	October 25, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240134887 A1	Apr. 25, 2024
US 20240232233 A9	Jul. 11, 2024

Publication Classification

Int. Cl.: G06F16/00 (20190101); G06F16/22 (20190101); G06F16/2458 (20190101); G06F16/28 (20190101)

U.S. Cl.:

CPC G06F16/288 (20190101); G06F16/2228 (20190101); G06F16/2477 (20190101);

Field of Classification Search

CPC: G06F (16/288); G06F (16/2228)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11922942	12/2023	Kilinc	N/A	G10L 15/26
2009/0192968	12/2008	Tunstall-Pedoe	706/53	G06N 5/02
2010/0198503	12/2009	Beckner	707/E17.001	G06F 16/29
2011/0307435	12/2010	Overell	706/46	G06F 16/367
2023/0004583	12/2022	Marmanis	N/A	G06F 16/287

OTHER PUBLICATIONS

“Data and Analytics Essentials: Data Fabric”, Gartner [Online]. Retrieved from the Internet: <URL: <https://www.gartner.com/en/documents/4003523>>, (Jul. 13, 2021), 6 pgs. cited by applicant

Bach, Nguyen, et al., “A Review of Relation Extraction”, Literature review for Language and Statistics, (2007), 15 pgs. cited by applicant

Panetta, Kasey, “5 Trends Appear on the Gartner Hype Cycle for Emerging Technologies, 2019”, Gartner, Inc., [Online]. Retrieved from the Internet: <URL: <https://www.gartner.com/smarterwithgartner/5-trends-appear-on-the-gartner-hype-cycle-for-emerging-technologies-2019>>, (Aug. 29, 2019), 6 pgs. cited by applicant

Paulheim, Heiko, “Knowledge Graph Refinement: A Survey of Approaches and Evaluation Methods”, Semantic Web, 8(3), IOS Press, (2016), 489-508. cited by applicant

Pawar, Sachin, et al., “Relation Extraction: A Survey”, (2017), 51 pgs. cited by applicant

Primary Examiner: Boccio; Vincent F

Attorney, Agent or Firm: SCHWEGMAN LUNDBERG & WOESSNER, P.A.

Background/Summary

BACKGROUND

(1) Knowledge graphs are commonly used to represent human-usable information in a computing system. For example, a knowledge graph may be used to store enterprise information in a business or other enterprise in a way that is manageable by a computing system. Enterprise users may retrieve information from the knowledge graph using the computing system.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) The present disclosure is illustrated by way of example and not limitation in the following figures.
- (2) FIG. 1 is a diagram showing one example of an environment for generating and/or maintaining a knowledge graph data structure.
- (3) FIG. 2 is a diagram showing a fragment of a knowledge graph and illustrating an example triple data unit including metadata describing a relationship represented by the example knowledge graph.
- (4) FIG. 3 is a flowchart showing one example of a process flow that may be executed in the environment of FIG. 1.
- (5) FIG. 4 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to incorporate extracted relationship data into the knowledge graph data structure.
- (6) FIG. 5 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system to generate and/or to supplement a knowledge graph data structure.
- (7) FIG. 6 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to respond to queries against the knowledge graph data structure from enterprise users, such as the enterprise user.
- (8) FIG. 7 is a diagram showing one example user interface page that may be provided to the user to facilitate a user query against the knowledge graph data structure.
- (9) FIG. 8 is a diagram showing another example of the user interface page of FIG. 7.
- (10) FIG. 9 is a diagram showing another example of a user interface page that may be provided to the user.
- (11) FIG. 10 shows another example of the user interface page of FIG. 9 including a feedback window.
- (12) FIG. 11 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to respond to a time-dependent query against the knowledge graph data structure.
- (13) FIG. 12 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to perform single-period processing of a time-dependent query.
- (14) FIG. 13 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to perform closed-closed processing of a time-dependent query against the knowledge graph data structure.
- (15) FIG. 14 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to perform closed-open processing of a time-dependent query against the knowledge graph data structure.
- (16) FIG. 15 is a diagram showing one example of a probability distribution based on confirmation points for a test relationship.
- (17) FIG. 16 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to perform open-closed processing of a time-dependent query against the knowledge graph data structure.
- (18) FIG. 17 is a diagram showing one example of a probability distribution based on confirmation points for a test relationship.
- (19) FIG. 18 is a flowchart showing one example of a process flow that may be performed by the knowledge graph computing system of FIG. 1 to perform open-open processing of a time-dependent query against the knowledge graph data structure.
- (20) FIG. 19 is a flowchart showing one example of a process flow that may be executed by the

knowledge graph computing system of FIG. 1 to execute multi-period processing of a time-dependent query against the knowledge graph data structure.

(21) FIG. 20 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to incorporate a time-dependent relationship into the knowledge graph data structure.

(22) FIG. 21 is a flowchart showing one example of a process flow that may be executed by the knowledge graph computing system of FIG. 1 to evaluate the veracity of time-dependent relationships in the knowledge graph data structure.

(23) FIG. 22 is a diagram showing one example of an environment implementing a knowledge graph data structure service.

(24) FIG. 23 is a diagram showing one example of a cloud environment for implementing a knowledge graph data structure service.

(25) FIG. 24 is a block diagram showing one example of an architecture for a computing device.

(26) FIG. 25 is a block diagram of a machine in the example form of a computer system within which instructions may be executed for causing the machine to perform any one or more of the methodologies discussed herein.

DETAILED DESCRIPTION

(27) Knowledge graphs can be used in various different computing systems to store and retrieve information about relationships. Consider an example in which a computing system uses a knowledge graph to store relationships that indicate skills held by enterprise employees. In this example, a user of the computing system may be a manager who has been tasked with implementing a machine learning algorithm. The user however may not be an expert in machine learning algorithms and may not be directly aware of any experts in machine learning algorithms within the enterprise employing the user. The user may make a query against the knowledge graph requesting the identity of employees who are skilled in machine learning algorithms. The knowledge graph computing system may return relationships stored by the knowledge graph data structure that indicate employees skilled in machine learning algorithms.

(28) Consider an example relationship “employee A is skilled in machine learning algorithms.” In this example, the subject entity is “employee A,” the object entity is “machine learning algorithms” and the predicate is “is skilled in.” The computing system may return the indicated relationship to the user, along with other relationships stored by the knowledge graph also indicating skill in machine learning algorithms. In this way, the user may identify enterprise employees who may be suitable to implement the machine learning algorithm that has been tasked to the user.

(29) Although a knowledge graph can be very useful in this scenario, the use of a knowledge graph in computing systems can present other challenges. For example, building a knowledge graph may be a time-consuming, and often manual, process. Also, the existence of relationships in a knowledge graph may not provide a user with an indication of the veracity of the relationship. Consider again the example above. In some examples, employee A may be skilled in machine learning algorithms that were state-of-the-art at the time that the relationship was added to the knowledge graph data structure but may not be skilled in more recent machine learning algorithms. In this case, the usefulness of the relationship provided to the user may be compromised. Also, in some examples, employee A may have been in a developer role at the time the relationship was added to the knowledge graph data structure but may have subsequently transitioned to a management role or another role that does not involve active coding, or may have left the enterprise. In these and other examples, the relationship “employee A is skilled in machine learning algorithms” may not be useful to the user.

(30) Various examples address these and other challenges by utilizing a knowledge graph data structure and associated systems and methods, as described herein. For example, a knowledge graph data structure may be constructed with triple data units. A triple data unit may describe a relationship. The triple data unit may comprise a subject subunit indicating a subject entity, and

object subunit indicating an object, and a predicate subunit describing a predicate of the subject entity and the object. A triple data unit may also include other subunits such as, for example a confidence subunit comprising data indicating a confidence in the accuracy of the relationship. Other example subunits of a triple data unit may include, for example, a confidentiality subunit indicating a level of confidentiality of the relationships stored by the triple data unit, an author subunit indicating an author of a document or documents from which the relationship was extracted, a location subunit indicating a location of the document or documents from which the relationship was extracted, and/or a last edit subunit indicating the last time that the document or documents from which the relationship was extracted have been edited.

(31) In some examples, a triple data unit may also include a time subunit storing data describing temporal properties of the relationship. The time subunit may include data describing the veracity of the relationship represented by the triple data unit at one or more times. For example, the time subunit may include an indication of one or more confirmation points-in-time. A confirmation point-in-time may be a point in time at which the relationship represented by the triple data unit is true. In some examples, the time subunit may also indicate one or more periods during which the relationship represented by the triple data unit is, or was, true.

(32) In various examples, a knowledge graph data structure may be generated from enterprise and/or external documents in an automated manner. Automated knowledge graph data structure generation may be based on an ontology and/or using a machine learning model. For example, an ontology-based technique and/or machine learning technique may be used to generate a knowledge graph data structure. Users may utilize a knowledge graph computing system to make queries against the knowledge graph data structure. The knowledge graph computing system may generate responses to the queries. Query responses provided to users may include descriptions of the relationships that are responsive to the query and a feedback request. The feedback request may prompt the user to provide feedback data describing a confidence value in one or more of the returned relationships. The feedback data may indicate, for example, whether a relationship is true or false and, in some examples when the relationship is, or was, true or false. In some examples, the feedback may also indicate a degree of confidence in the truth of the indicated relationship.

(33) The knowledge graph computing system may utilize the feedback to generate labeled training data. For example, labeled training data may include an indication of the relationship (e.g., subject entity, object, and predicate) as well as an indication of a confidence value, where the confidence value indicates a likelihood that the relationship is true. Training data generated in this way by the knowledge graph computing system may be used to train or re-train a machine learning model for extracting additional relationships from documents. In this way, the generation and supplementing of the knowledge graph data structure may be performed by the knowledge graph computing system in an automated or semiautomated manner.

(34) FIG. 1 is a diagram showing one example of an environment **100** for generating and/or maintaining a knowledge graph data structure **104**. The knowledge graph data structure **104** may be maintained by a knowledge graph computing system **102**. The knowledge graph computing system **102** may be or include any suitable computing device or devices, such as one or more servers. The knowledge graph computing system **102**, for example, may be implemented in an on-premise and/or cloud-based computing environment.

(35) The knowledge graph computing system **102** may maintain the knowledge graph data structure **104** at one or more machine-readable media, which may be implemented in one or more databases or other suitable data stores. The knowledge graph data structure **104** may include various triple data units **142**, **144**. Each triple data unit **142**, **144** may describe a relationship and may include various subunits including, for example, a subject subunit, and object subunit, and a predicate subunit. In the example of FIG. 1, the knowledge graph data structure **104** describes a knowledge graph comprising entities **132**, **134**, predicates **138**, **140**, and a literal **136**. It will be appreciated that the knowledge graph data structure **104** may describe additional entities,

predicates, and literals. The components shown in FIG. 1 are illustrated for example purposes.

(36) The triple data unit **144** describes a relationship in which the entity **132** is the subject, the entity **134** is the object, and the predicate **138** is the predicate. Accordingly, the triple data unit **144** may include a subject subunit describing the entity **132**, a predicate subunit describing the predicate **138**, and an object subunit describing the entity **134**. The triple data unit **142** describes a relationship in which the entity **134** is the subject, the predicate **140** is the predicate, and the literal **136** is the object. In some examples, a literal may be an object but may not be a subject of a relationship. Accordingly, the triple data unit **142** may include a subject subunit describing the entity **134**, a predicate subunit describing the predicate **140**, and an object subunit describing the literal **136**. The example triple data units **144**, **142** may also include other data describing the indicated relationships such as, for example, confidence data describing the veracity of the relationships, confidentiality level data describing a confidentiality of the relationships, and so on, as described herein. Additional examples describing the structure of triple data units are provided herein including, for example, with respect to FIG. 2.

(37) The knowledge graph computing system **102** may implement various subsystems **112**, **114**, **116**, **118** to facilitate training, modification, and the use of the knowledge graph data structure **104**. Subsystems **112**, **114**, **116**, **118** may include executable code, computing hardware, such as one or more hardware processors, one or more machine-readable media, and/or the like.

(38) The access layer subsystem **112** may interface with a user **130** via a user computing device **128**. The user **130** may be any user with privileges to query the knowledge graph data structure. The user **130** may be an internal or enterprise user associated with an enterprise implementing the knowledge graph computing system **102**. In some examples, the user **130** may be an external user outside of the enterprise implementing the knowledge graph computing system **102**. FIG. 1 also shows an administrative user **124** who may access the knowledge graph computing system **102** via a user computing device **126**. The user computing devices **126**, **128** may be any suitable computing device such as, for example, a laptop computer, a desktop computer, a tablet computer, a mobile computing device, and/or the like. Although two users **124**, **130** are shown in FIG. 1, it will be appreciated that multiple different users may access the knowledge graph computing system **102**, for example, to utilize and/or administer the knowledge graph data structure **104** as described herein.

(39) The access layer subsystem **112** may support queries against the knowledge graph data structure **104**. For example, the user **130** may generate a query using the user computing device **128** and provide the query to the knowledge graph computing system **102** via the access layer subsystem **112**. The access layer subsystem **112** may execute the query against the knowledge graph data structure **104** and return one or more relationships described by the knowledge graph data structure **104** that are responsive to the query. In some examples, the access layer subsystem **112** provides the user **130** with query results via an interface page that also includes a request for feedback data about one or more of the indicated relationships.

(40) The environment **100** also includes a query assistant system **120**. The query assistant system **120** is shown separate from the knowledge graph computing system **102**. In some examples, however, the query assistant system **120** may be implemented as a subsystem of the knowledge graph computing system **102**. The query assistant system **120** may be programmed to assist the user **130** in generating queries against the knowledge graph data structure **104**. For example, the user **130** may direct a natural language question to the query assistant system **120**. The query assistant system **120** may be programmed to convert the natural language question to one or more queries against the knowledge graph data structure **104**. In some examples, the query assistant system **120** may make queries against the knowledge graph data structure **104** in the course of assisting the user **130**. For example, when the user **130** makes a natural language query to the query assistant system **120**, the query assistant system **120** may respond by making one or more queries against the knowledge graph data structure **104**. The query assistant system **120** may return results of the query

to the user **130** and/or may use results of the query to recommend a query against the knowledge graph data structure **104** to the user **130**.

(41) The confidence subsystem **114** may be configured to manage confidence values for one or more relationships indicated by triple data units **142**, **144** of the knowledge graph data structure **104**. In some examples, the confidence subsystem **114** may generate an initial confidence value for a relationship when the relationship is added to the knowledge graph data structure **104** as a triple data unit **142**, **144**. The confidence subsystem **114** may also update the confidence value associated with a relationship in response to feedback data provided by enterprise users such as the user **130**.

(42) The sense dictionary subsystem **116** may be configured to resolve sense disambiguation in human-readable text. The sense dictionary subsystem **116** may be used, for example, when a relationship is written to the knowledge graph data structure **104**, when modifying relationships described by triple data units at the knowledge graph data structure **104**, in responding to queries against the knowledge graph data structure **104**, and/or the like.

(43) Consider the term “bank.” This may refer to the edge of a body of water or a financial institution. The sense dictionary subsystem **116** may receive an indication of a relationship including a subject entity, predicate, and object. The sense dictionary subsystem **116** may return an indication of a sentence associated with at least one of the subject entity, the predicate, or the object. Returning to the example above, the relationship: “XYZ is a bank” may have a subject entity “XYZ,” a predicate “is a,” and an object “bank.” The sense dictionary subsystem **116** may utilize the relationship and, in some examples, text from which the relationship was extracted and provide sense data indicating that the object “bank” refers to a financial institution.

(44) The ontology subsystem **118** may be configured to store and/or administer an ontology. The ontology may describe relationships between words and corresponding entities, predicates, and/or literals that may be described by triple data units at the knowledge graph data structure **104**. In some examples, the ontology subsystem **118** may communicate with the relationship extractor system **108**, as described herein, to facilitate the extraction of relationships from documents **110**, as described herein.

(45) The ontology may comprise information that may be represented in the knowledge graph data structure **104**. For example, the ontology may include information about relationships of different types including, for example, what sorts of entities can be subjects of the relationship, what sorts of objects and/or literals can be objects of the relationship, and the like. In some examples, the ontology may store words that may be used, for example, in natural language to represent a relationship. Consider an example relationship with a predicate “works on,” such as, for example, A works on B. One example implementation of this relationship may be “Employee X works on New_Project.” The ontology may list words indicating relationships of the relevant type such as “works on,” “is assigned to,” “is on the project team for,” and/or the like. In this example, all of the following phrases may indicate the relationship: “Employee X works on New_Project.” “Employee X is assigned to New_Project.” “Employee X is on the project team for New_Project.”

(46) In some examples the ontology managed by the ontology subsystem **118** may also store information about time dependencies of relationship. For example, a relationship may be time dependent or time independent. A relationship that is time dependent may have a veracity that changes at different points-in-time. Consider again the example relationship “works on.” This relationship may be true at some times and false at other times and, therefore, may be indicated at the ontology to be time dependent. Other examples of relationships that may be time dependent include “lives,” “works for,” “has temperature,” and “has market price.”

(47) A relationship may be time independent if its veracity is always the same for all time. Consider another example relationship “born on.” This relationship may be time independent because the birth date of an entity (a person) does not change with time and therefore a statement of a birth date is either true or false for all time. In some examples, the ontology subsystem **118** supports the additional time dependency categories such as, for example, time-dependency-ignored. A

relationship may be classified as time-dependency-ignored, for example, if the relationship can change with time but practically does not change very often. Consider the following example relationships: “has currency,” and “is the capital of” Although it is possible for a country or similar entity to change its currency and it is possible for a country or similar entity to change its capital, such changes do not occur very often. As such it may not be practical to manage these relationships as time dependent in the knowledge graph data structure **104**.

(48) In some examples, the ontology managed by the ontology subsystem **118** may indicate a type of time dependency for time-dependent relationships. For example, a relationship may have a single period of time dependency or a multi-period time dependency. A relationship having a single-period time dependency may be true only during a single continuous time and false at all other times. A relationship having a multi-period time dependency may be true at multiple, non-continuous time periods. Consider the example relationship “lives.” An entity (a human) only lives once. Therefore, a relationship indicating that a particular human being lives is true only during one continuous time period and the relationship “lives” has a single-period time dependency. Consider another example relationship “works for.” An entity, such as a human, may work for the same company during two different non-continuous time periods. Accordingly, the relationship “works for” as a multi-period time dependency.

(49) In some examples, the ontology managed by ontology subsystem **118** may support other types of time dependencies such as, for example, a multi-period-ignored time dependency. A relationship may have a multi-period-ignored time dependency if, for example, it is technically possible, but very unlikely, for the relationship to be true during two non-continuous time periods. Consider the example relationship “is Chancellor of Germany.” This relationship may technically have a multi-period time dependency because it is possible to be Chancellor of Germany during more than one time. Practically, however, it may be very rare for an individual to be Chancellor of Germany during multiple noncontinuous time periods.

(50) The knowledge graph computing system **102** also includes a relationship extractor system **108** and a training database **106**. The relationship extractor system **108** may comprise executable code, computing hardware, such as one or more hardware processors, one or more machine-readable media, and/or the like. Although the relationship extractor system **108** is shown as a component of the knowledge graph computing system **102**, in various examples, the relationship extractor system **108** may be implemented as a separate computing system.

(51) The relationship extractor system **108** may be configured to extract relationships from documents **110**. The documents **110** may be or include any suitable document associated with the enterprise implementing the knowledge graph computing system **102**. The documents **110** may include, for example, internal documents generated by the enterprise as well as external documents generated by parties outside the enterprise. Examples of document types that may be included in the documents **110** include product documents, intranet pages, transcriptions, Internet pages, books, email messages, and/or the like. Product documents may include documents describing a product or service provided by the enterprise implementing the knowledge graph computing system **102**. Intranet pages may include pages from an internal network accessible to some or all users associated with the enterprise implementing the knowledge graph computing system **102**.

Transcriptions may include documents that reflect an automated and/or manual transcription of meetings, conversations, and/or other interactions between people associated with the enterprise.

(52) In some examples, the relationship extractor system **108** operates utilizing the ontology subsystem **118** and/or the ontology managed thereby. The ontology may include a set of words describing various different entities, literals, and predicates. The relationship extractor system **108** may examine the documents **110** to find instances of words indicating a combination of a subject entity, a predicate, and an object entity or littoral.

(53) In other examples, the relationship extractor system **108** is programmed to implement a machine learning model. The machine learning model may be, for example, a classifier that is

trained to extract relationships from the documents **110**. The machine learning model may be trained using labeled training data stored at the training database **106**. In some examples, as described herein, the training database **106** may be populated with training data that is derived from feedback data received from enterprise users such as the user **130**.

(54) In some examples, the relationship extractor system **108** is programmed to extract relationships from documents **110** using the ontology managed by the ontology subsystem **118** and a machine learning model. For example, the relationship extractor system **108** may initially extract relationships from the documents **110** using the ontology subsystem **118** as described herein. This may generate and/or supplement the knowledge graph data structure **104**. As users, such as the user **130**, query the knowledge graph data structure **104** and provide feedback data, the feedback data may be used to generate labeled training data that may be stored at the training database **106**. A machine learning model may be trained using the training data. When the machine learning model is trained, the relationship extractor system **108** may begin to extract relationships from the documents **110** using the machine learning model.

(55) When the relationship extractor system **108** determines a relationship from the documents **110**, it may add the relationship to the knowledge graph data structure **104**. This may include generating a new triple data unit corresponding to the relationship, and/or supplementing an existing triple data unit describing the relationship. For example, a relationship extracted by the relationship extractor system **108** may already be represented by a triple data unit at the knowledge graph data structure **104**. If this is the case, the relationship extractor system **108** may, rather than writing a new triple data unit, modify the existing triple data unit describing the relationship. For example, the relationship extractor system **108** may update a confidence value for the relationship stored at a confidence subunit of an existing triple data unit based on the document **110** from which the relationship was extracted.

(56) In some examples, the relationship extractor system **108** may provide an extracted relationship to the administrative user **124**, who may access the extracted relationship via the user computing device **126**. The administrative user **124** may provide an initial confidence value for the relationship. The confidence value provided by the administrative user **124** may become the confidence value for the relationship and/or may be supplemented with other data such as, for example, the document **110** from which the relationship was extracted.

(57) FIG. 2 is a diagram showing a fragment **200** of a knowledge graph and illustrating an example triple data unit **201** describing a relationship **203** represented by the example knowledge graph fragment **200**. The knowledge graph fragment **200** includes entities **202**, **204**, **206**, **208** and predicates **210**, **212**, **214**. In the example of FIG. 2, entities **202**, **204**, **206**, **208** are indicated by circles. Predicates **210**, **212**, **214** are indicated by arrows.

(58) The entity **202** is “ODM (Operational Data Model)”. The entity **204** is “data model.” In the example of FIG. 2, the entity **204** data model has been assigned a sense “semantic.” For example, the sense “semantic” may differentiate the entity **204** from entities for database data models, physical data models, and other types of data models. In some examples, the sense may have been assigned by the sense dictionary subsystem **116**. The predicate **210** is “is a.” Accordingly, the relationship **203** can be stated as “ODM is a data model.”

(59) The knowledge graph fragment **200** shows other example relationships. For example, the relationship “data model is expressed by entity relationship model” is shown by entity **204**, predicate **212**, and entity **206**. An additional example relationship “power designer uses entity relationship model” is expressed by entity **208**, predicate **214**, and entity **206**.

(60) FIG. 2 shows an example triple data unit **201** describing the relationship **203**. The triple data unit **201** comprises various subunits including, for example, the subject subunit **220**, a predicate subunit **222**, and an object subunit **224**. The subject subunit **220**, the predicate subunit **222**, and the object subunit **224** describe the subject, predicate, and object of the relationship **203**.

(61) The triple data unit **201** may comprise other subunits describing other data about the

relationship **203**. A confidence subunit **226** may store a confidence value describing a level of confidence in the veracity of the relationship **203**. A confidentiality subunit **228** may store an indication of a confidentiality level associated with the relationship **203**.

(62) A time subunit **232** may store data describing temporal properties of the relationship **203** represented by the triple data unit **201**. For example, the time subunit **232** may store one or more confirmation points-in-time for the relationship **203**. A confirmation point-in-time may be a point-in-time at which the relationship **203** is known to be true. Confirmation points-in-time may be compiled and stored at the time subunit **232** in any suitable manner. In some examples, confirmation points-in-time may be derived from documents **110**. For example, if a document **110** indicates that the relationship **203** is true, then the date of the document (e.g., the publication date, the last edit date, and/or the like) may be a confirmation point-in-time.

(63) Also, in some examples, the knowledge graph computing system **102** may derive a confirmation point-in-time for one relationship from one or more other relationships. For example, the ontology managed by the ontology subsystem **118** may indicate relationships that are implied by other relationships. Consider the example relationship “lived” with respect to a human subject entity. Numerous other relationships may imply the truth of the relationship “lived” at a particular point-in-time. Examples of such relationships may include “visited,” “worked for,” “is Chancellor of,” and/or the like.

(64) In some examples, the confirmation points-in-time stored at the time subunit **232** may include one or more start dates and/or end dates for the relationship **203**. For example, the ontology managed by the ontology subsystem **118** may indicate events that are start dates and/or end dates for different relationships. Consider again the example relationship “lived” with respect to an individual. The ontology may indicate that a relationship “was born on” stored at the knowledge graph data structure **104** and having the same subject entity may indicate a start date for the relationship “lived.” Similarly, the ontology may indicate that a relationship “died on” stored at the knowledge graph data structure **104** and having the same subject entity may indicate an end date for the relationship “lived.”

(65) Various additional subunits are shown with respect to source data **230**. The source data **230** comprises information describing one or more documents **110** from which the relationship **203** was extracted. In this example, the source data **230** comprises a name, author, file location, and last edit date for two documents **110** from which the relationship **203** was extracted. The source data **230**, in this example, also comprises an indication of the sentence or other text unit from the respective documents **110** that was used to derive the relationship **203**.

(66) FIG. 3 is a flowchart showing one example of a process flow **300** that may be executed in the environment **100** of FIG. 1. For example, the process flow **300** may be executed, in whole or in part, by the knowledge graph computing system **102**. At operation **302**, the knowledge graph computing system **102** (e.g., the relationship extractor system **108** thereof) may access a document **110**. At operation **304**, the knowledge graph computing system **102** (e.g., the relationship extractor system **108** thereof) may extract a relationship from the document **110**. The relationship may be extracted from the document **110** in any suitable manner. In some examples, the relationship extractor system **108** accesses the ontology subsystem **118** and retrieves entries from the ontology corresponding to words and/or phrases in the document **110**. The relationship extractor system **108** may utilize the ontology entries to determine relationships represented by the language of the document **110**. In other examples, the relationship extractor system **108** may implement a trained machine learning model. The relationship extractor system **108** may apply the trained machine learning model to the document **110** to return one or more relationships indicated by the text of the document **110**.

(67) At operation **306**, the knowledge graph computing system **102** (e.g., the relationship extractor system **108** thereof) may incorporate the extracted relationship into the knowledge graph data structure **104**. This may include, for example, writing a new triple data unit at the knowledge graph

data structure **104** and/or modifying an existing triple data unit at the knowledge graph data structure **104**. The operations **302**, **304**, **306** are described with respect to a single document **110**. It will be appreciated, however, that operations **302**, **304**, **306** may be performed for multiple documents **110** either in series or in parallel.

(68) At operation **308**, the knowledge graph computing system **102** (e.g., the access layer subsystem **112** thereof) may provide data to a user **130** in response to a query against the knowledge graph data structure **104** made by the user **130**. The access layer subsystem **112** may provide the results, for example, by serving a user interface page to the user **130** via the user computing device **128**. Example user interface pages that may be used to provide query results to the user **130** are described herein with respect to FIGS. **8-10**. The user interface page may include a description of the relationship or relationships from the knowledge graph data structure **104** that are responsive to the query. In some examples, the user interface page may also solicit feedback data from the user **130**, where the feedback data describes a veracity of the relationship or relationships returned in response to the query.

(69) The knowledge graph computing system **102** (e.g., the access layer subsystem **112** thereof) receives feedback data from the user **130** at operation **310**. The feedback data may indicate the user's opinion of the veracity of the relationship or relationships provided. For example, if a relationship is true, the user may indicate that the relationship is true. If the relationship is not true or otherwise not useful to the user, this may also be indicated by the feedback data. In some examples, a user may indicate that a relationship is true, wrong, outdated, or does not fit the query made by the user.

(70) The knowledge graph computing system **102** may utilize the feedback data to modify a confidence value for the returned relationship at operation **312**. For example, as described herein, a triple data unit at the knowledge graph data structure **104** describing the relationship may include a confidence subunit storing a confidence value for the relationship. The knowledge graph computing system **102** may update the confidence value based on the feedback data and write the updated confidence value data to the confidence subunit of the appropriate triple data unit.

(71) The confidence value for the relationship may be updated in various different ways. For example, if a user indicates that a relationship is false or wrong, the knowledge graph computing system **102** may decrease the confidence value for the relationship and write the decreased confidence value to the confidence subunit of the corresponding triple data unit. Equation [1] below shows one example way that the confidence value for a relationship may be decreased upon receiving negative user feedback:

$$c.sub.t = c.sub.t - t.sub.u * d \quad [1]$$

In Equation [1], $c.sub.t$ is the trust value for the relationship; $t.sub.u$ is a trust score associated with the user providing the feedback data; and d is the maximum decrease in the confidence value that is permitted to occur based on feedback data from a single user.

(72) If a user indicates that a relationship is true, the knowledge graph computing system may increase the confidence value for the relationship and write the increased confidence value of the confidence subunit of the corresponding triple data unit. Equation [2] below shows one example way that the confidence value for a relationship may be increased upon receiving positive user feedback:

$$c.sub.t = c.sub.t + t.sub.u * i \quad [2]$$

In Equation [1], $c.sub.t$ is the trust value for the relationship; $t.sub.u$ is a trust score associated with the user providing the feedback data; and i is the maximum increase in the confidence value that is permitted to occur based on feedback data from a single user.

(73) The trust score $t.sub.u$ for a user, such as the user **130** or the administrative user **124**, may be determined in various different ways. In some examples the trust score is initially determined based on a role of the user **124**, **130** within the enterprise. In some examples, the trust score for a user may be updated based on relationships added to the knowledge graph data structure **104** by that

user. For example, if a user **124**, **130** adds a relationship that is determined to be false, the user's trust score may be decreased. On the other hand, if the user adds relationships that are determined to be true, the user's trust score may be increased.

(74) FIG. **4** is a flowchart showing one example of a process flow **400** that may be executed by the knowledge graph computing system **102** of FIG. **1** to incorporate extracted relationship data into the knowledge graph data structure **104**. The process flow **400** provides one example way of executing the operation **306** of the process flow **300**.

(75) At operation **402**, the knowledge graph computing system **102** accesses relationship data describing a relationship. The relationship data may be accessed, for example, from the relationship extractor system **108** which may have extracted the relationship from a document **110**. The relationship data may indicate a subject entity, a predicate, and an object of the relationship, which may be an entity or a literal.

(76) At operation **404**, the knowledge graph computing system **102** may determine if any of the elements of the relationship (e.g., the subject entity, the predicate, and/or the object) are expressed in language that could be interpreted in multiple senses. If the relationship does include one or more elements that can be expressed in multiple senses, the knowledge graph computing system **102** may utilize the sense dictionary subsystem **116** to assign a proper sense to the one or more elements of the relationship at operation **406**. For example, the knowledge graph computing system **102** and/or the sense dictionary subsystem **116** may examine context words and/or phrases from the document **110**. The context words and/or phrases may be matched with senses indicated at the sense dictionary subsystem **116**. The determined sense for a subject entity may be written to the subject subunit of the corresponding triple data unit. The determined sense for a predicate may be written to the predicate subunit of the corresponding triple data unit. The determined sense for an object may be written to the object subunit of the corresponding triple data unit. It will be appreciated that one, two, or all three of the subject entity, the predicate, and the object may be assigned a sense.

(77) At operation **408**, the knowledge graph computing system **102** may determine if the extracted relationship is a new relationship or if it has already been described at the knowledge graph data structure **104** with an existing triple data unit. If the relationship has not yet been represented at the knowledge graph data structure **104**, the knowledge graph computing system **102** may write a new triple data unit for the relationship to the knowledge graph data structure **104** at operation **410**. The knowledge graph computing system **102** may populate the triple data unit and its various subunits with data describing the relationship. For example, the subject entity may be described by the subject subunit, the object may be described at the object subunit, and the predicate may be described at the predicate subunit.

(78) In some examples, one or more entities of the newly extracted relationship may also be part of other relationships represented at the knowledge graph data structure **104** by other triple data units. For example, the subject entity of the extracted relationship may be a subject entity and/or an object of one or more other relationships. Similarly, if the object of the extracted relationship is an entity, then that object entity may be a subject or an object of one or more other relationships. When one or more entities of the newly extracted relationship are part of other relationships represented at the knowledge graph data structure **104** by other triple data units, knowledge graph computing system **102** may note the other relationships, for example, at the triple data units and/or at a separate storage location.

(79) In some examples, the knowledge graph computing system **102** may also populate other subunits of the triple data unit. For example, the knowledge graph computing system **102** may determine an initial confidentiality level for the relationship. The initial confidentiality level may be determined in any suitable manner. In some examples, the ontology managed by the ontology subsystem **118** may indicate a confidentiality level of different types of relationships. The knowledge graph computing system **102** may consult the ontology subsystem **118** to determine the

appropriate confidentiality level for the relationship, based on an indicated confidentiality level for relationships of the same type. In some examples, the confidentiality value is based on the document **110** from which the relationship was extracted. For example, the confidentiality of the relationship may be set equal to the confidentiality of the document **110** from which the relationship was extracted. The confidentiality value determined for the relationship may be written to a confidentiality subunit of the triple data unit.

(80) In some examples, the knowledge graph computing system **102** (e.g., the confidence subsystem **114** thereof) determines the confidence value for the relationship. The confidence value for the extracted relationship may be determined in various different ways. In some examples, the extracted relationship may be verified by the administrative user **124** as described herein. In these examples, the confidence value for the extracted relationship may be set equal to a trust value associated with the administrative user **124** as given by Equation [3] below in which $c_{sub.t}$ is the confidence value for the extracted relationship and $t_{sub.u}$ is the trust level of the administrative user:

$$c_{sub.t} = t_{sub.u} \quad [3]$$

In some examples, a relationship may be added to the knowledge graph data structure **104** manually by the administrative user **124**. In these examples, the confidence value for the relationship may also be determined according to Equation [3].

(81) In some examples, the knowledge graph computing system **102** (e.g., the confidence subsystem **114** thereof) determines the confidence value for the relationship without considering an administrative user **124**, for example, when the administrative user **124** does not verify an extracted relationship. For example, the confidence value may be determined based on various factors including a confidence value for the relationship extractor system **108** and a confidence value associated with an author of the document **110** from which the relationship was extracted. Equation [4] provides an example of how these factors may be considered:

$$(82) \quad c_t = \frac{(c_a + c_r)}{2 \times p} \quad [4]$$

In Equation [4], $c_{sub.a}$ is a confidence value or trust score associated with an author of the document **110** from which the relationship was extracted. Also, $c_{sub.r}$ is a confidence value associated with the relationship extractor system **108**. This may indicate an accuracy of the algorithm being used by the relationship extractor system **108** to extract relationships from documents **110**. In some examples, the relationship extractor system **108** may have different confidence values based on the algorithm used to extract the relationship. For example, when the relationship extractor system **108** is using an ontology from the ontology subsystem **118** to extract relationships, it may have one confidence value. When it is using a trained machine learning algorithm, it may have another different confidence value. The value p represents an optional punishing value. For example, the larger the punishing value p , the lower the confidence value determined for automatically extracted relationships.

(83) Referring back to operation **408**, if the extracted relationship is not a new relationship, then the knowledge graph computing system **102** may modify, at operation **412**, an existing triple data unit at the knowledge graph data structure **104** based on the extracted relationship. For example, if the extracted relationship is not a new relationship, there may be an existing triple data unit at the knowledge graph data structure **104** representing the extracted relationship. The knowledge graph computing system, at operation **412**, may update the existing triple data unit. For example, the knowledge graph computing system **102** may supplement the existing triple data unit to include data describing the document **110** from which the relationship was extracted. This may lead to a triple data unit indicating multiple documents **110**, such as, for example, the example described with respect to FIG. 2.

(84) The knowledge graph computing system **102** may also, in some examples, update the confidence value associated with the relationship at the existing triple data unit. For example, the

knowledge graph computing system **102** may re-determine a confidence value for the relationship in view of the new extraction of the relationship from the document **110**. For example, the knowledge graph computing system **102** may consider a confidence value for the relationship extractor system **108**, confidence value or trust level of an author of the document **110**, and or other suitable factors.

(85) FIG. 5 is a flowchart showing one example of a process flow **500** that may be executed by the knowledge graph computing system **102** to generate and/or to supplement a knowledge graph data structure **104**. At operation **502**, the knowledge graph computing system **102** (the relationship extractor system **108** thereof) may access a document **110**. At operation **504**, the knowledge graph computing system **102** may extract a relationship from the document **110** using an ontology, such as an ontology managed by the ontology subsystem **118**. The ontology may include words or phrases indicating relationships. The relationship extractor system **108** may find words and/or phrases in the document **110** that match the words or phrases in the ontology and extract relationships therefrom based on the ontology. Consider the following example. The ontology may indicate that the phrase “is a” represents a predicate linking a subject entity referenced prior to the predicate phrase and an object referenced after the predicate phrase. When the relationship extractor system **108** encounters said example text “Berlin is a city,” it may apply the ontology to extract a relationship in which the subject entity is “Berlin,” the predicate is “is a,” and the object is “city.”

(86) At operation **506**, the knowledge graph computing system **102** may incorporate the extracted relationship into the knowledge graph data structure **104**, for example, as described herein with respect to the process flow **400**. At operation **508**, the knowledge graph computing system **102** may receive feedback data describing the relationship extracted at operation **504**. For example, an enterprise user, such as the user **130**, may make a query against the knowledge graph data structure **104** that returns the extracted relationship. The enterprise user may provide feedback data describing the relationship. The feedback data may indicate that the relationship is true, or may indicate that the relationship is not true, outdated, or otherwise not useful. If the enterprise user indicates that the relationship is true, the confidence value for the relationship may be updated at its corresponding triple data unit, e.g., the confidence subunit thereof. If the enterprise user indicates that the relationship is not true, outdated, or otherwise not useful, the confidence value for the relationship may be reduced.

(87) At operation **510**, the knowledge graph computing system **102** may update the training database **106** based on the feedback data received at operation **508**. This may include modifying one or more labels at the training database **106**. For example, the training database **106** may indicate the extracted relationship and may be labeled with various data including, for example, text from a document or documents **110** from which the relationship was extracted, and an indication of the confidence value for the relationship. Updating the training database **106** may include, for example, updating the confidence value associated with the relationship at training database **106**. At operation **512**, the knowledge graph computing system **102** may train and/or retrain a machine learning model to extract relationships from documents **110**. When the machine learning model is trained, it may be used by the relationship extractor system **108** to extract additional relationships from documents **110**.

(88) FIG. 6 is a flowchart showing one example of a process flow **600** that may be executed by the knowledge graph computing system **102** (e.g., the access layer subsystem **112** thereof) to respond to queries against the knowledge graph data structure **104** from enterprise users, such as the user **130**. At operation **602**, the access layer subsystem **112** may receive a query against the knowledge graph data structure **104**. The query may have originated from the user **130** and may be received via the user computing device **128**.

(89) At operation **604**, the access layer subsystem **112** may identify one or more triple data units at the knowledge graph data structure **104** that are responsive to the query. At operation **606**, the

access layer subsystem **112** may determine a confidentiality level match between the user **130** and any relationships that are responsive to the query. For example, enterprise users may have a confidentiality level indicating a level of confidential information that the user is permitted to see. A relationship stored at a triple data unit in the knowledge graph data structure (indicated “TDU” in FIG. **6**) may also have a confidentiality level stored at a confidence subunit of the triple data unit. Determining the confidentiality level match may include determining whether the confidentiality level of the user **130** is sufficient to permit the user to view a returned relationship. If the user's confidentiality level does not meet the confidentiality level indicated by the triple data unit of a responsive relationship, that relationship may not be shown to the user.

(90) At operation **608**, the access layer subsystem **112** may serve a user interface page to the user **130** via the user computing device **128**. The user interface page may include a representation of one or more relationships responsive to the user's query. The interface page may also include a prompt, such as, for example, a visual and/or a textual prompt. The prompt may solicit feedback data from the user.

(91) FIG. **7** is a diagram showing one example user interface page **702** that may be provided to the user **130** to facilitate a user query against the knowledge graph data structure **104**. The user interface page **702** may be provided to the user **130** via the user computing device **128**. The user interface page **702** comprises a query field **701** with a text field **704** where the user **130** may enter a query or question to be executed against the knowledge graph data structure **104**. The user **130** may initiate the query by selecting a search button **706**.

(92) The user interface page **702** also includes a result field **703**. For example, a version of the user interface page **702** comprising the query field **701** only may be provided to the user to prompt the user to provide a query. In some examples, the full user interface page **702** including the query field **701** and the result field **703** may be provided with results of the user's query. The result field **703** indicates a relationship that is responsive to the user's query at a relationship field **710**. The relationship indicated that the result field may include a subject entity, a predicate, and an object. The subject entity, predicate, and object may be read from the triple data unit associated with the relationship at the knowledge graph data structure **104**. A metadata field **708** includes various metadata regarding the relationship indicated that the relationship field **710**. Metadata at the metadata field **708** may also have been read from the triple data unit associated with the relationship at the knowledge graph data structure **104**. In this example, the metadata includes sentences indicating the relationship from various documents **110** from which the relationship is extracted along with additional information about the documents including the document name, the document author, a location of the documents, a confidentiality level associated with the documents, and a last edit date for the documents.

(93) The user interface page **702** also includes a Context Graph button **712**. The user **130** may select the Context Graph button **712** to replace the metadata field **708** with the graphical representation of a portion of the knowledge graph including the relationship indicated in the relationship field **710**. FIG. **8** is a diagram showing another example of the user interface page **702**. The user interface page **702** may appear as depicted in FIG. **8**, for example, when the user **130** selects the context graph button **712**. In FIG. **8**, the user interface page **702** includes the context graph field **802**. The context graph field **802** graphically depicts a portion of the knowledge graph including the relationship indicated at relationship field **710**. In some examples, the user may select a Back To Detail Page button **804** to return to the version of the user interface page **702** depicted in FIG. **7**.

(94) FIG. **9** is a diagram showing another example of a user interface page **902** that may be provided to the user **130**. The user interface page **902** may be provided to the user **130** via the user computing device **128**. The user interface page **902** comprises a query field **904** where the user **130** may enter a query or question to be executed against the knowledge graph data structure **104**. The user **130** may initiate the query by selecting a Search button **906**. The user interface page **902** also

includes a result field **908**. For example, a version of the user interface page **902** comprising the query field **904** may be provided to the user **130** only to prompt the user **130** to provide a query. In some examples, the full user interface page **902** including the query field **904** and the result field **908** may be provided with results of the user's query.

(95) The result field **908** includes indications of three relationships that are potentially responsive to the user query indicated at the query field **904**. Each relationship is indicated by a row at the result field **908**. Additional entries in the row provide additional information about each of the potentially responsive relationships such as, for example, a Source Category, a Likelihood, a link that may be selected to provide Further Details about the relationship, and a Voting field soliciting feedback data. The relationships shown at the result field **908** may have been read from the knowledge graph data structure **104**. For example, each relationship may be represented at the knowledge graph data structure **104** by a triple data unit, as described herein.

(96) The Source Category associated with a relationship may indicate a source of the document **110** from which the relationship is extracted. In some examples, the Source Category may provide an indication of the confidentiality level associated with the relationship. For example, a Source Category of “internal” may only be available to users who are internal to the enterprise implementing the knowledge graph data structure **104**. A Source Category indicating “external” may be available to users outside the enterprise as well.

(97) The Likelihood value for a relationship may be an indication of the confidence value for the relationship, as described herein. In the example of FIG. **9**, the Likelihood or confidence value is provided as a percentage. It will be appreciated, however, that a likelihood or confidence value may be provided in various other formats in addition to or instead of as a percent.

(98) The Voting field, as shown, includes buttons that may be selected by the user to provide an indication of whether each indicated relationship is positive or negative. For example, the user **130** may select the buttons showing a check to indicate that the corresponding relationship is true and responsive to the user's query. The user **130** may select the button and ask to indicate that the corresponding relationship is either false or not responsive to the user's query.

(99) FIG. **10** shows another example of the user interface page **902** including a feedback window **1002**. In some examples, the feedback window **1002** may be shown to the user **130** when the user selects the negative feedback button from the voting field of a relationship. The feedback window **1002** prompts user **130** to indicate whether the relationship does not fit the user's query, is outdated, or is wrong. In the example of FIG. **10**, the feedback window **1002** also includes a field where the user may enter an “other” statement indicating why a given relationship is negative. As described herein, the knowledge graph computing system **102** may modify the confidence value of a relationship based on the feedback data that may be provided such as, for example, as shown in FIGS. **9** and **10**.

(100) In some examples, a user, such as the user **130**, may make a time-dependent query against the knowledge graph data structure **104**. A time-dependent query may be a query having a result that depends on the veracity of at least one test relationship at a test point-in-time. Consider the following example query: “Was Elizabeth II Queen of the United Kingdom in 1967?” This example query against the knowledge graph data structure **104** may be a time-dependent query testing whether the relationship “Elizabeth II is Queen of the United Kingdom” is true at a test point-in-time, which in this example is 1967. Time-dependent queries may also be used, as described herein, for various other purposes including, for example, to verify relationships before the relationships are added to the knowledge graph data structure **104**, and/or to verify relationships that have already been included in the knowledge graph data structure **104**.

(101) In some examples, the knowledge graph computing system **102** is configured to respond to time-dependent queries using the knowledge graph data structure **104**. FIG. **11** is a flowchart showing one example of a process flow **1100** that may be executed by the knowledge graph computing system **102** to respond to a time-dependent query against the knowledge graph data

structure **104**.

(102) At operation **1102**, the knowledge graph computing system **102** may receive a time-dependent query against the knowledge graph data structure **104**. The time-dependent query may indicate a test relationship and a test point-in-time. At operation **1104**, the knowledge graph computing system **102** may determine if the test relationship is a time-dependent relationship. This may include consulting the ontology managed by the ontology subsystem **118**, for example, by querying the ontology subsystem **118**. If the test relationship is not time-dependent, the knowledge graph computing system **102** may return an error in response to the time-dependent query at operation **1106**.

(103) If the test relationship is time-dependent, then the knowledge graph computing system **102** may determine, at operation **1108**, whether the test relationship is single-period time-dependent (indicated as S_P in FIG. **11**) or multi-period time-dependent (indicated at M_P in FIG. **11**). This may be performed, for example, using the ontology. For example, the ontology may indicate whether it is single-period time-dependent and or multi-period time-dependent. If the test relationship is single-period time-dependent, then the knowledge graph computing system **102** may execute single-period processing of the query at operation **1110**. Examples describing single-period processing of the time-dependent query are described herein with respect to FIGS. **12-18**. If the test relationship is multi-period time-dependent, then the knowledge graph computing system **102** may execute multi-period processing of the query at operation **1112**. Examples describing multi-period processing of the time-dependent query are described herein with respect to FIG. **19**.

(104) FIG. **12** is a flowchart showing one example of a process flow **1200** that may be executed by the knowledge graph computing system **102** to perform single-period processing of a time-dependent query. The process flow **1200** is one example way that the knowledge graph computing system **102** may perform operation **1110** of the process flow **1100**.

(105) The process flow **1200** may process time-dependent queries with single-period test relationships based on whether the boundaries of the single period are open or closed. A boundary may be closed if the knowledge graph data structure **104** includes a date indicating a date for the boundary. Consider again the example relationship “lived.” If the test relationship includes the predicate “lived,” then a beginning boundary of the test relationship may be closed if the knowledge graph data structure **104** indicates a relationship “was born on” having the same subject entity as the test relationship. An end boundary for the test relationship may be closed if the knowledge graph data structure **104** indicates a relationship “died in” having the same subject entity as the test relationship. In some examples, a test relationship may be closed-closed, open-closed, closed-open, or open-open. If the test relationship is closed-closed, both the beginning and the end of the single period are closed or defined. If the test relationship is closed-open, then a beginning boundary of the single period may be closed or defined while an end boundary of the single period may not be defined. If the test relationship is open-closed, then an end boundary of the single period may be closed or defined while the beginning boundary of the single period may not be defined. If the test relationship is open-open, then neither boundary of the single period may be defined.

(106) At operation **1202**, the knowledge graph computing system **102** may determine if the test relationship is closed-closed. If the test relationship is closed-closed, then the knowledge graph computing system **102** may perform closed-closed processing of the query at operation **1204**, for example, as described herein with respect to FIG. **13**.

(107) If the test relationship is not closed-closed, then the knowledge graph computing system **102** may, at operation **1206**, utilize one or more confirmation points-in-time associated with the test relationship to generate a time span distribution describing the test relationship. This may include, for example, fitting a probability distribution to the confirmation points-in-time. The confirmation points-in-time are points-in-time at which the test relationship is known or believed to be true. The confirmation points may be stored, for example, at a time subunit of the triple data unit associated

with the test relationship at the knowledge graph data structure **104**. In some examples, the knowledge graph computing system **102** accesses the confirmation points-in-time from other triple data units at the knowledge graph data structure **104**. For example, the ontology managed by the ontology subsystem **118** may describe other relationships that, if based on the same subject entity and/or object, indicate the truth or falsity of the test relationship. The knowledge graph computing system **102** may execute one or more queries against the knowledge graph data structure **104** to locate triple data units indicating relationships that provide confirmation points-in-time for the test relationship.

(108) The distribution may be any suitable probability distribution such as, for example, a normal distribution, a Gaussian distribution, and/or the like. The knowledge graph computing system **102** may generate the probability distribution based on the relationships indicated at the knowledge graph data structure **104**. For example, the mean of the probability distribution may be determined based on the mean date/time of the confirmation points-in-time. The variance of the probability distribution may also be determined from the confirmation points-and-time. In some examples, the confirmation points-in-time may be drawn from knowledge graph data structures stored by other tenants at a cloud environment, as described herein below. This may increase the number of confirmation points-in-time used to generate the probability distribution and, thereby, may increase the accuracy of the probability distribution.

(109) At operation **1208**, the knowledge graph computing system **102** may determine if the test relationship is closed-open. If the test relationship is closed-open, then at operation **1210** the knowledge graph computing system **102** may perform closed-open processing of the query using the time span distribution determined at **1206**. An example of closed-open processing is provided herein with respect to FIG. **14**. If the test relationship is not closed-open, then the knowledge graph computing system may, at operation **1212**, determine whether the test relationship is open-closed. If the test relationship is open-closed, then the knowledge graph computing system **102** may perform open-closed processing of the query using the time span distribution at operation **1214**. An example of open-closed processing is provided herein with respect to FIG. **16**. If the test relationship is not open-closed, then the test relationship may be open-open. The knowledge graph computing system **102** may perform open-open processing of the query using the time span distribution at operation **1216**. An example of open-open processing is provided herein with respect to FIG. **18**.

(110) FIG. **13** is a flowchart showing one example of a process flow **1300** that may be executed by the knowledge graph computing system **102** to perform closed-closed processing of a time-dependent query against the knowledge graph data structure **104**. At operation **1302**, the knowledge graph computing system **102** may determine whether the test point-in-time (indicated as “P-I-T” in FIG. **13**) is within the closed time period of the test relationship. For example, the knowledge graph computing system **102** may determine whether the test point-in-time is after a defined beginning of the time period and before a defined end of the time period. If the test point-in-time is in the closed time period of the relationship, then the knowledge graph computing system **102** may, at operation **1304**, return an indication that the test relationship is true at the test point-in-time. If the test point-in-time is not in the closed time period of the test relationship, then the knowledge graph computing system **102** may, at operation **1306**, return an indication that the test relationship is false at the test point-in-time. The indication that the test relationship is true or false at the test point-in-time may make up all or a portion of a response to the time-dependent query. The knowledge graph computing system **102** may send the response to the time-dependent query to the computing device that made the time-dependent query, such as, for example, a user computing device of a user, the query assistant system **120**, and/or the like.

(111) FIG. **14** is a flowchart showing one example of a process flow **1400** that may be executed by the knowledge graph computing system **102** to perform closed-open processing of a time-dependent query against the knowledge graph data structure **104**. In this example, a beginning boundary of the time period of the test relationship is defined or closed, but the end boundary is not

defined or closed. The process flow **1400** shows one example manner in which the operation **1210** of process flow **1200** may be executed.

(112) At operation **1402**, the knowledge graph computing system **102** may count a number of confirmation points-in-time for the test relationship that occur after the test point-in-time. For example, confirmation points-in-time for the test relationship occurring after the test point-in-time may tend to indicate that the test relationship is true at the test point-in-time. If the number of confirmation points-in-time for the test relationship occurring after the test point-in-time is greater than the threshold at operation **1404**, then the knowledge graph computing system **102** may, at operation **1406**, return an indication that the test relationship is true at the test point-in-time.

(113) If the number of confirmation points-in-time for the test relationship occurring after the test point-in-time is not greater than the threshold, then the knowledge graph computing system **102** may, at operation **1408**, return a value of the probability distribution of confirmation points-in-time as a veracity of the test relationship at the test point-in-time.

(114) FIG. **15** is a diagram **1500** showing one example of a probability distribution **1502** based on confirmation points for a test relationship. The probability distribution **1502** is shown relative to a horizontal axis **1503** and a vertical axis **1501**. The horizontal axis **1503** indicates time. The vertical axis **1501** indicates a probability that the test relationship is true at the time indicated by the horizontal axis **1503**. As described herein, the probability distribution **1502** may be generated based on the confirmation points-in-time for the test relationship and may be based on any suitable distribution.

(115) The diagram **1500** illustrates one example way that the operation **1408** of the process flow **1400** may be performed. For example, the test relationship represented by the probability distribution **1502** is closed-open. Accordingly, the distribution is positioned such that the probability exceeds zero at a start date indicated by “startDate.” The test point-in-time is indicated on the horizontal axis **1503** by “testpointInTime.” Accordingly, in the example of FIG. **15**, returning the value of the probability distribution **1502** at the test point-in-time may include returning the value indicated by “p.”

(116) FIG. **16** is a flowchart showing one example of a process flow **1600** that may be executed by the knowledge graph computing system **102** to perform open-closed processing of a time-dependent query against the knowledge graph data structure **104**. In this example, an end boundary of the time period of the test relationship is defined or closed, but the beginning boundary is not defined or closed. The process flow **1600** shows one example manner in which the operation **1214** of process flow **1200** may be executed.

(117) At operation **1602**, the knowledge graph computing system **102** may count a number of confirmation points-in-time for the test relationship that occur before the test point-in-time. For example, confirmation points-in-time for the test relationship occurring before the test point-in-time may tend to indicate that the test relationship is true at the test point-in-time. If the number of confirmation points-in-time for the test relationship occurring before the test point-in-time is greater than the threshold at operation **1604**, then the knowledge graph computing system **102** may, at operation **1606**, return an indication that the test relationship is true at the test point-in-time.

(118) If the number of confirmation points-in-time for the test relationship occurring before the test point-in-time is not greater than the threshold, then the knowledge graph computing system **102** may, at operation **1608**, return a value of the probability distribution of confirmation points-in-time as a veracity of the test relationship at the test point-in-time.

(119) FIG. **17** is a diagram **1700** showing one example of a probability distribution **1702** based on confirmation points for a test relationship. The probability distribution **1702** is shown relative to a horizontal axis **1703** and a vertical axis **1701**. The horizontal axis **1703** indicates time. The vertical axis **1701** indicates a probability that the test relationship is true at the time indicated by the horizontal axis **1703**. As described herein, the probability distribution **1702** may be generated based on the confirmation points-in-time for the test relationship and may be based on any suitable

distribution.

(120) The diagram **1700** illustrates one example way that the operation **1608** of the process flow **1600** may be performed. For example, the test relationship represented by the probability distribution **1702** is open-closed. Accordingly, the distribution is position such that the probability drops to zero at an end date indicated by “endDate.” The test point-in-time is indicated on the horizontal axis **1703** by “testpointInTime.” Accordingly, in the example of FIG. 17, returning the value of the probability distribution **1702** at the test point-in-time may include returning the value indicated by “p.”

(121) FIG. 18 is a flowchart showing one example of a process flow **1800** that may be performed by the knowledge graph computing system **102** to perform open-open processing of a time-dependent query against the knowledge graph data structure **104**. For example, the process flow **1800** shows one example manner in which the operation **1216** of the process flow **1200** may be performed. In this example, neither an end boundary nor a beginning boundary of the time period of the test relationship is defined or closed. Instead, both are open.

(122) At operation **1802**, the knowledge graph computing system **102** may determine if there are any confirmation points-in-time for the test relationship. For example, the knowledge graph computing system **102** may consult a time subunit of the triple data unit describing the test relationship to determine if it indicates any confirmation points-in-time. If there are no confirmation points-in-time for the test relationship, the knowledge graph computing system **102** may, at operation **1804**, return an uncertainty base value as the probability that the test relationship is true at the test point-in-time. The uncertainty base value may be a value indicating that the knowledge graph data structure **104** does not indicate a likelihood that the test relationship is true at the test point-in-time. In some examples, the uncertainty base value may be 0.5 or 50%.

(123) If there are one or more confirmation points-in-time for the test relationship, the knowledge graph computing system **102**, at operation **1806**, may center the probability distribution determined at operation **1206**. This may include positioning a mean of the probability distribution at an average time of the confirmation points-in-time for the test relationship. If there is only a single confirmation point-in-time for the test relationship, this may involve positioning the mean of the probability distribution at the confirmation point-in-time. If there are multiple confirmation points-in-time for the test relationship, this may include positioning the mean of the probability distribution at an average time of the confirmation points-in-time. At operation **1808**, the knowledge graph computing system **102** may return the value of the centered probability distribution at the test point-in-time as the probability that the test relationship is true at the test point-in-time.

(124) FIG. 19 is a flowchart showing one example of a process flow **1900** that may be executed by the knowledge graph computing system **102** to execute multi-period processing of a time-dependent query against the knowledge graph data structure **104**. For example, the process flow **1900** shows one example manner of executing the operation **1112** of the process flow **1100**.

(125) At operation **1902**, the knowledge graph computing system **102** may determine a set of confirmation points-in-time that are closest to the test point-in-time. In some examples, this includes selecting a single confirmation point-in-time that is the closest after the test point-in-time and a single confirmation point-in-time that is the closest before the test point-in-time. Also, in some examples, this may include selecting the closest n points-in-time after the test point-in-time and the closest n points-in-time before the test point-in-time.

(126) At operation **1904**, the knowledge graph computing system **102** may increment a base probability by the probability of each of the closest confirmation points-in-time selected at operation **1902**. The base probability may be any suitable probability such as, for example, 0.5 or 50%. In some examples, the operation **1904** may be performed according to the Equation [5] below:

$$(127) \ S = \sum_{i=1}^n \text{Math. } v + p_i \quad [5]$$

In the Equation [5], the value S is the veracity of the test relationship at the test point-in-time and may be the result of the incrementing; the value n is the number of closest confirmation points-in-time selected at operation **1902**; the value v is the base probability; and the value p.sub.i is the value of the probability distribution at each of the closest confirmation points-in-time. At operation **1906**, the knowledge graph computing system **102** may return the incremented probability value (e.g., S) as the probability that the test relationship is true at the test point-in-time.

(128) In some examples, the techniques described herein for responding to time-dependent queries may be utilized to verify the veracity of a time-dependent relationship as the relationship is incorporated into the knowledge graph data structure **104**. A time-dependent relationship may be a relationship that is provided with a time-dependent statement. Consider the example statement from above: “Elizabeth II was Queen of the United Kingdom in 1967.” This relationship depends on a test point-in-time, which is 1967. Accordingly, before a time-dependent relationship is written to the knowledge graph data structure **104** as a triple data unit, it may be desirable to verify the time-dependent relationship. In some examples, a triple data unit may indicate a time-dependent relationship as an additional predicate or object of the relationship according to a suitable format such as, for example, RDF*. Referring to the example statement above the information “in 1967” could be added to the triple data unit indicating “Elizabeth II was Queen of the United Kingdom.”

(129) FIG. **20** is a flowchart showing one example of a process flow **2000** that may be executed by the knowledge graph computing system **102** to incorporate a time-dependent relationship into the knowledge graph data structure **104**. At operation **2002**, the knowledge graph computing system **102** may access a candidate time-dependent relationship. The candidate time-dependent relationship may have been extracted from a document **110**, received from and/or verified by a user **124**, **130**. At operation **2004**, the knowledge graph computing system **102** may determine a time-dependent probability of the candidate time-dependent relationship. This may be performed, for example, as described herein with respect to FIGS. **11-19**. At operation **2006**, the knowledge graph computing system **102** determines if the probability determined that operation **2004** is less than a threshold confidence value. If the probability is not less than the threshold confidence value, then at operation **2008** the knowledge graph computing system may incorporate the candidate time-dependent relationship into the knowledge graph data structure **104**, for example, as described herein.

(130) If the returned probability is less than the threshold confidence value, the knowledge graph computing system may, at operation **2010**, obtain a confirmation of the candidate time-dependent relationship. This may include, for example, prompting a user, such as the administrative user **124**, to verify the veracity of the candidate time-dependent relationship. If the user confirms the candidate time-dependent relationship at operation **2012**, then the knowledge graph imputing system may incorporate the time-dependent relationship into the knowledge graph data structure **104** at operation **2008**. If the user does not confirm the veracity of the candidate time-dependent relationship at operation **2012**, then the knowledge graph computing system **102** may perform exception processing at operation **2014**. Exception processing may include, for example, postponing or canceling the addition of the time-dependent relationship to the knowledge graph data structure **104**.

(131) In some examples, the knowledge graph computing system **102** may analyze the knowledge graph data structure **104** to evaluate the veracity of time-dependent relationships reflected by the knowledge graph data structure **104**. For example, one or more relationships described by triple data units, as described herein, may be based on a point-in-time. FIG. **21** is a flowchart showing one example of a process flow **2100** that may be executed by the knowledge graph computing system **102** to evaluate the veracity of time-dependent relationships in the knowledge graph data structure **104**.

(132) At operation **2102**, the knowledge graph computing system **102** may access the knowledge graph data structure **104**. At operation **2104**, the knowledge graph computing system **102** may determine a probability of one or more time-dependent relationships represented by triple data units at the knowledge graph data structure **104**. This may be performed, for example, as described herein with respect to FIGS. **11-19**.

(133) At operation **2106**, the knowledge graph computing system **102** may determine if the knowledge graph data structure **104** includes triple data units describing any relationships that have a probability that is less than a threshold confidence value. If no time-dependent relationships have a probability less than the threshold confidence value, then the process flow **2100** may conclude at operation **2108**. If any time-dependent relationships have a probability less than the threshold confidence value, then the knowledge graph computing system **102** may, at operation **2110**, prompt a user, such as the administrative user **130**, to confirm the time-dependent relationships having a probability less than the threshold confidence value.

(134) At operation **2112**, the knowledge graph computing system **102** may determine if all of the time-dependent relationships having a probability less than the threshold confidence value have been confirmed. If all have been confirmed, then the process flow **2100** may conclude at operation **2108**. If any time-dependent relationships having a probability less than the threshold confidence value are not confirmed at operation **2112**, then the knowledge graph computing system **102** may perform exception processing at operation **2114**. Exception processing may include, for example, removing from the knowledge graph data structure **104** triple data units describing time-dependent relationships having a probability less than the threshold confidence value and being unconfirmed. In another example, exception processing may include submitting time-dependent relationships having a probability less than the threshold confidence value and being unconfirmed for additional review.

(135) FIG. **22** is a diagram showing one example of an environment **2200** implementing a knowledge graph data structure service **2202**. The knowledge graph data structure service **2202** may operate in a manner similar to the knowledge graph computing system **102** described herein. In some examples the knowledge graph data structure service **2202** and a knowledge graph data structure **2204** may be executed at a cloud deployment **2201**. The cloud deployment **2201** may include one or more data centers implementing one or more virtual and/or hardware servers. The cloud deployment **2201** executes the knowledge graph data structure service **2202** and implements the knowledge graph data structure **2204** at one or more data stores.

(136) A user **2230** may access the knowledge graph data structure service **2202** via a user computing device **2228** as described herein. In some examples, the knowledge graph data structure service **2202** may be executed in conjunction with an enterprise application **2222**. In the example of FIG. **22**, the enterprise application **2222** is executed at an enterprise site **2220**. The enterprise site **2220** may be an on-premise and/or cloud-implemented computing system for executing the enterprise application **2222**. The user **2230**, in some examples, may access the knowledge graph data structure service **2202** via the enterprise application **2222**. For example, the user **2230** may utilize the enterprise application **2222** to make queries of the knowledge graph data structure **2204** via the knowledge graph data structure service **2202**. Replies to queries made by the enterprise application **2222** may be utilized in the processing of the enterprise application **2222**, for example, so as to serve the user **2230**.

(137) FIG. **22** also shows an administrative user **2224**. The administrative user **2224** may access the enterprise site **2220**, the enterprise application **2222**, and/or the knowledge graph data structure service **2202** via a user computing device **2226**. In some examples, the administrative user **2224** may perform functions similar to those of the administrative user **124** described herein. The administrative user **2224**, and some examples, may have additional administrative roles with respect to the knowledge graph data structure service **2202** and/or the enterprise application **2222**.

(138) In some examples, time-dependent queries and relationships may be validated in a cloud

environment using confirmation points-in-time shared between tenants of the cloud environment. FIG. 23 is a diagram showing one example of a cloud environment **2300** for implementing a knowledge graph data structure service **2202**. The cloud environment **2300** includes a cloud deployment **2301**, which may be similar to the cloud deployment **2201** of FIG. 22. In this example, the cloud deployment **2301** comprises a plurality of different tenancies. Each tenancy may be associated with different customers such as, for example, different customer enterprises.

(139) In the example of FIG. 23, the cloud deployment comprises a knowledge graph data structure service **2202** and separate tenant knowledge graph data structures **2302**, **2304**, **2306**. Each tenant knowledge graph data structure **2302**, **2304**, **2306** may be associated with a separate tenancy at the cloud deployment **2301**. Also, each of the respective tenant knowledge graph data structures **2302**, **2304**, **2306** may include various triple data units indicating relationships.

(140) The relationships stored by a different tenant knowledge graph data structures **2302**, **2304**, **2306**, in some examples, are proprietary to the respective tenants. Accordingly, it may not be desirable to give one tenant access to another tenant's knowledge graph data structure **2302**, **2304**, **2306**. In some examples, however, the knowledge graph data structure service **2202** may be configured to identify equivalent relationships from among the tenant knowledge graph data structures **2302**, **2304**, **2306** and to share some or all confirmation points-in-time for the equivalent relationships. This may involve, for example, writing to the time subunits of triple data units corresponding to the equivalent relationship's confirmation points-in-time from the other tenant knowledge graph data structures **2302**, **2304**, **2306**.

(141) In some examples, confirmation points-in-time may be shared between tenant knowledge graph data structures **2302**, **2304**, **2306** based on a confidentiality level of the document **110** from which the confirmation points-in-time are extracted. For example, if a confirmation point-in-time in the tenant knowledge graph data structure **2302** was extracted from a document **110** that was internal to the respective tenant, that confirmation point-in-time may not be shared to the tenant knowledge graph data structures **2304**, **2306** of the other tenants. On the other hand, if the tenant knowledge graph data structure **2302** includes a triple data unit reflecting a confirmation point-in-time extracted from a non-confidential document, then that confirmation point-in-time may be shared to triple data units indicating equivalent relationships at the other tenant knowledge graph data structures **2304**, **2306**.

EXAMPLES

(142) Example 1 is a computing system for utilizing relationship data, the computing system comprising: at least one processor programmed to perform operations comprising: receiving, from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship, and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining, using the first plurality of confirmation points-in-time, that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure; determining a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query to the second computing system.

(143) In Example 2, the subject matter of Example 1 optionally includes the operations further comprising: determining that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and determining that a number of the first plurality of confirmation points-in-time are before the first test point-in-time, the response to the first time-dependent query indicating that the first test relationship is true at the first test point-in-time.

(144) In Example 3, the subject matter of any one or more of Examples 1-2 optionally include the operations further comprising: determining that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and fitting a probability distribution to the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

(145) In Example 4, the subject matter of any one or more of Examples 1-3 optionally include the operations further comprising: determining that the beginning of the first time period is not defined by the knowledge graph data structure; determining that the end of the first time period is not defined by the knowledge graph data structure; and positioning a probability distribution using the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

(146) In Example 5, the subject matter of any one or more of Examples 1-4 optionally include the operations further comprising: determining that the first test relationship is a multi-period relationship; selecting a first confirmation point-in-time of the first plurality of confirmation points-in-time that is after the first test point-in-time; and selecting a second confirmation point-in-time of the first plurality of confirmation points-in-time that is before the first test point-in-time, the veracity of the first test relationship at the first test point-in-time being based at least in part on a first probability distribution centered at the first confirmation point-in-time and a second probability distribution centered on the second confirmation point-in-time.

(147) In Example 6, the subject matter of any one or more of Examples 1-5 optionally include the accessing of the confirmation data comprising accessing a first triple data unit from the knowledge graph data structure, the first triple data unit corresponding to the first test relationship.

(148) In Example 7, the subject matter of any one or more of Examples 1-6 optionally include the accessing of the confirmation data comprising accessing a plurality of triple data units from the knowledge graph data structure.

(149) In Example 8, the subject matter of any one or more of Examples 1-7 optionally include the operations further comprising: accessing a candidate relationship comprising a candidate subject entity, a candidate predicate, a candidate object, and a candidate point-in-time; accessing a second plurality of confirmation points-in-time at which the candidate relationship is true; determining veracity of the candidate relationship at the candidate point-in-time using the second plurality of confirmation points-in-time; determining that the veracity of the candidate relationship at the candidate point-in-time is greater than a threshold value; and after determining that the veracity of the candidate relationship at the candidate point-in-time is greater than the threshold value, incorporating the candidate relationship into the knowledge graph data structure.

(150) In Example 9, the subject matter of any one or more of Examples 1-8 optionally include the operations further comprising: accessing a second test triple data unit from the knowledge graph data structure, the second test triple data unit describing a second test relationship and a second test point-in-time; accessing second confirmation data from the knowledge graph data structure, the second confirmation data describing a second plurality of confirmation points-in-time at which the second test relationship is true; determining that a veracity of the second test relationship at the second test point-in-time is less than a threshold value; and obtaining a confirmation of the second test relationship.

(151) In Example 10, the subject matter of any one or more of Examples 1-9 optionally include the knowledge graph data structure being stored at a first tenancy of a plurality of tenancies at a cloud deployment, the operations further comprising: accessing a second knowledge graph data structure stored at a second tenancy of the plurality of tenancies at the cloud deployment; and using the second knowledge graph data structure to determine at least a portion of the first plurality of confirmation points-in-time at which the first test relationship is true.

(152) Example 11 is a method for utilizing relationship data in a knowledge graph computing

system, the method comprising: receiving, by the knowledge graph computing system and from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing, by the knowledge graph computing system, confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining, by the knowledge graph computing system and using the first plurality of confirmation points-in-time, that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure; determining, by the knowledge graph computing system, a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query, by the knowledge graph computing system and to the second computing system.

(153) In Example 12, the subject matter of Example 11 optionally includes determining, by the knowledge graph computing system, that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and determining, by the knowledge graph computing system, that a number of the first plurality of confirmation points-in-time are before the first test point-in-time, the response to the first time-dependent query indicating that the first test relationship is true at the first test point-in-time.

(154) In Example 13, the subject matter of any one or more of Examples 11-12 optionally include determining, by the knowledge graph computing system, that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and fitting a probability distribution to the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

(155) In Example 14, the subject matter of any one or more of Examples 11-13 optionally include determining, by the knowledge graph computing system, that the beginning of the first time period is not defined by the knowledge graph data structure; determining, by the knowledge graph computing system, that the end of the first time period is not defined by the knowledge graph data structure; and positioning a probability distribution using the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

(156) In Example 15, the subject matter of any one or more of Examples 11-14 optionally include determining, by the knowledge graph computing system, that the first test relationship is a multi-period relationship; selecting, by the knowledge graph computing system, a first confirmation point-in-time of the first plurality of confirmation points-in-time that is after the first test point-in-time; and selecting, by the knowledge graph computing system, a second confirmation point-in-time of the first plurality of confirmation points-in-time that is before the first test point-in-time, the veracity of the first test relationship at the first test point-in-time being based at least in part on a first probability distribution centered at the first confirmation point-in-time and a second probability distribution centered on the second confirmation point-in-time.

(157) In Example 16, the subject matter of any one or more of Examples 11-15 optionally include the accessing of the confirmation data comprising accessing a first triple data unit from the knowledge graph data structure, the first triple data unit corresponding to the first test relationship.

(158) In Example 17, the subject matter of any one or more of Examples 11-16 optionally include accessing, by the knowledge graph computing system, a candidate relationship comprising a candidate subject entity, a candidate predicate, a candidate object, and a candidate point-in-time; accessing, by the knowledge graph computing system, a second plurality of confirmation points-in-time at which the candidate relationship is true; determining, by the knowledge graph computing

system, a veracity of the candidate relationship at the candidate point-in-time using the second plurality of confirmation points-in-time; determining, by the knowledge graph computing system, that the veracity of the candidate relationship at the candidate point-in-time is greater than a threshold value; and after determining that the veracity of the candidate relationship at the candidate point-in-time is greater than the threshold value, incorporating the candidate relationship into the knowledge graph data structure.

(159) In Example 18, the subject matter of any one or more of Examples 11-17 optionally include accessing, by the knowledge graph computing system, a second test triple data unit from the knowledge graph data structure, the second test triple data unit describing a second test relationship and a second test point-in-time; accessing, by the knowledge graph computing system, second confirmation data from the knowledge graph data structure, the second confirmation data describing a second plurality of confirmation points-in-time at which the second test relationship is true; determining, by the knowledge graph computing system, that a veracity of the second test relationship at the second test point-in-time is less than a threshold value; and obtaining, by the knowledge graph computing system, a confirmation of the second test relationship.

(160) In Example 19, the subject matter of any one or more of Examples 11-18 optionally include the knowledge graph data structure being stored at a first tenancy of a plurality of tenancies at a cloud deployment, the method further comprising: accessing, by the knowledge graph computing system, a second knowledge graph data structure stored at a second tenancy of the plurality of tenancies at the cloud deployment; and using, by the knowledge graph computing system, the second knowledge graph data structure to determine at least a portion of the first plurality of confirmation points-in-time at which the first test relationship is true.

(161) Example 20 is a machine-readable medium comprising instructions thereon that, when executed by at least one processor, cause the at least one processor to perform operations comprising: receiving, from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship, and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining, using the first plurality of confirmation points-in-time, that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure; determining a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query to the second computing system.

(162) FIG. 24 is a block diagram 2400 showing one example of a software architecture 2402 for a computing device. The architecture 2402 may be used in conjunction with various hardware architectures, for example, as described herein. FIG. 24 is merely a non-limiting example of a software architecture and many other architectures may be implemented to facilitate the functionality described herein. A representative hardware layer 2404 is illustrated and can represent, for example, any of the above referenced computing devices. In some examples, the hardware layer 2404 may be implemented according to the architecture of the computer system of FIG. 25.

(163) The representative hardware layer 2404 comprises one or more processing units 2406 having associated executable instructions 2408. Executable instructions 2408 represent the executable instructions of the software architecture 2402, including implementation of the methods, modules, subsystems, and components, and so forth described herein and may also include memory and/or storage modules 2410, which also have executable instructions 2408. Hardware layer 2404 may also comprise other hardware as indicated by other hardware 2412 which represents any other hardware of the hardware layer 2404, such as the other hardware illustrated as part of the

architecture **2402**.

(164) In the example architecture of FIG. **24**, the software architecture **2402** may be conceptualized as a stack of layers where each layer provides particular functionality. For example, the software architecture **2402** may include layers such as an operating system **2414**, libraries **2416**, middleware layer **2418**, applications **2420**, and presentation layer **2444**. Operationally, the applications **2420** and/or other components within the layers may invoke API calls **2424** through the software stack and access a response, returned values, and so forth illustrated as messages **2426** in response to the API calls **2424**. The layers illustrated are representative in nature and not all software architectures have all layers. For example, some mobile or special purpose operating systems may not provide a middleware layer **2418**, while others may provide such a layer. Other software architectures may include additional or different layers.

(165) The operating system **2414** may manage hardware resources and provide common services. The operating system **2414** may include, for example, a kernel **2428**, services **2430**, and drivers **2432**. The kernel **2428** may act as an abstraction layer between the hardware and the other software layers. For example, the kernel **2428** may be responsible for memory management, processor management (e.g., scheduling), component management, networking, security settings, and so on. The services **2430** may provide other common services for the other software layers. In some examples, the services **2430** include an interrupt service. The interrupt service may detect the receipt of an interrupt and, in response, cause the architecture **2402** to pause its current processing and execute an interrupt service routine (ISR) when an interrupt is accessed.

(166) The drivers **2432** may be responsible for controlling or interfacing with the underlying hardware. For instance, the drivers **2432** may include display drivers, camera drivers, Bluetooth® drivers, flash memory drivers, serial communication drivers (e.g., Universal Serial Bus (USB) drivers), Wi-Fi® drivers, NFC drivers, audio drivers, power management drivers, and so forth depending on the hardware configuration.

(167) The libraries **2416** may provide a common infrastructure that may be utilized by the applications **2420** and/or other components and/or layers. The libraries **2416** typically provide functionality that allows other software modules to perform tasks in an easier fashion than to interface directly with the underlying operating system **2414** functionality (e.g., kernel **2428**, services **2430** and/or drivers **2432**). The libraries **2416** may include system libraries **2434** (e.g., C standard library) that may provide functions such as memory allocation functions, string manipulation functions, mathematic functions, and the like. In addition, the libraries **2416** may include API libraries **2436** such as media libraries (e.g., libraries to support presentation and manipulation of various media format such as MPEG4, H.264, MP3, AAC, AMR, JPG, PNG), graphics libraries (e.g., an OpenGL framework that may be used to render 2D and 3D in a graphic content on a display), database libraries (e.g., SQLite that may provide various relational database functions), web libraries (e.g., WebKit that may provide web browsing functionality), and the like. The libraries **2416** may also include a wide variety of other libraries **2438** to provide many other APIs to the applications **2420** and other software components/modules.

(168) The middleware layer **2418** (also sometimes referred to as frameworks) may provide a higher-level common infrastructure that may be utilized by the applications **2420** and/or other software components/modules. For example, the middleware layer **2418** may provide various graphic user interface (GUI) functions, high-level resource management, high-level location services, and so forth. The middleware layer **2418** may provide a broad spectrum of other APIs that may be utilized by the applications **2420** and/or other software components/modules, some of which may be specific to a particular operating system or platform.

(169) The applications **2420** includes built-in applications **2440** and/or third-party applications **2442**. Examples of representative built-in applications **2440** may include, but are not limited to, a contacts application, a browser application, a book reader application, a location application, a media application, a messaging application, and/or a game application. Third-party applications

2442 may include any of the built-in applications **2440** as well as a broad assortment of other applications. In a specific example, the third-party application **2442** (e.g., an application developed using the Android™ or iOS™ software development kit (SDK) by an entity other than the vendor of the particular platform) may be mobile software running on a mobile operating system such as iOS™, Android™, Windows® Phone, or other mobile computing device operating systems. In this example, the third-party application **2442** may invoke the API calls **2424** provided by the mobile operating system such as operating system **2414** to facilitate functionality described herein.

(170) The applications **2420** may utilize built-in operating system functions (e.g., kernel **2428**, services **2430** and/or drivers **2432**), libraries (e.g., system libraries **2434**, API libraries **2436**, and other libraries **2438**), and middleware layer **2418** to create user interfaces to interact with users of the system. Alternatively, or additionally, in some systems interactions with a user may occur through a presentation layer, such as presentation layer **2444**. In these systems, the application/module “logic” can be separated from the aspects of the application/module that interact with a user.

(171) Some software architectures utilize virtual machines. In the example of FIG. **24**, this is illustrated by virtual machine **2448**. A virtual machine creates a software environment where applications/modules can execute as if they were executing on a hardware computing device. A virtual machine is hosted by a host operating system (operating system **2414**) and typically, although not always, has a virtual machine monitor **2446**, which manages the operation of the virtual machine as well as the interface with the host operating system (i.e., operating system **2414**). A software architecture executes within the virtual machine such as an operating system **2450**, libraries **2452**, frameworks/middleware **2454**, applications **2456** and/or presentation layer **2458**. These layers of software architecture executing within the virtual machine **2448** can be the same as corresponding layers previously described or may be different.

(172) Modules, Components and Logic

(173) Certain embodiments are described herein as including logic or a number of components, modules, or mechanisms. Modules may constitute either software modules (e.g., code embodied (1) on a non-transitory machine-readable medium or (2) in a transmission signal) or hardware-implemented modules. A hardware-implemented module is a tangible unit capable of performing certain operations and may be configured or arranged in a certain manner. In example embodiments, one or more computer systems (e.g., a standalone, client, or server computer system) or one or more hardware processors may be configured by software (e.g., an application or application portion) as a hardware-implemented module that operates to perform certain operations as described herein.

(174) In various embodiments, a hardware-implemented module may be implemented mechanically or electronically. For example, a hardware-implemented module may comprise dedicated circuitry or logic that is permanently configured (e.g., as a special-purpose processor, such as a field programmable gate array (FPGA) or an application-specific integrated circuit (ASIC)) to perform certain operations. A hardware-implemented module may also comprise programmable logic or circuitry (e.g., as encompassed within a general-purpose processor or another programmable processor) that is temporarily configured by software to perform certain operations. It will be appreciated that the decision to implement a hardware-implemented module mechanically, in dedicated and permanently configured circuitry, or in temporarily configured circuitry (e.g., configured by software) may be driven by cost and time considerations.

(175) Accordingly, the term “hardware-implemented module” should be understood to encompass a tangible entity, be that an entity that is physically constructed, permanently configured (e.g., hardwired), or temporarily or transitorily configured (e.g., programmed) to operate in a certain manner and/or to perform certain operations described herein. Considering embodiments in which hardware-implemented modules are temporarily configured (e.g., programmed), each of the hardware-implemented modules need not be configured or instantiated at any one instance in time.

For example, where the hardware-implemented modules comprise a general-purpose processor configured using software, the general-purpose processor may be configured as respective different hardware-implemented modules at different times. Software may accordingly configure a processor, for example, to constitute a particular hardware-implemented module at one instance of time and to constitute a different hardware-implemented module at a different instance of time.

(176) Hardware-implemented modules can provide information to, and receive information from, other hardware-implemented modules. Accordingly, the described hardware-implemented modules may be regarded as being communicatively coupled. Where multiple of such hardware-implemented modules exist contemporaneously, communications may be achieved through signal transmission (e.g., over appropriate circuits and buses that connect the hardware-implemented modules). In embodiments in which multiple hardware-implemented modules are configured or instantiated at different times, communications between such hardware-implemented modules may be achieved, for example, through the storage and retrieval of information in memory structures to which the multiple hardware-implemented modules have access. For example, one hardware-implemented module may perform an operation, and store the output of that operation in a memory device to which it is communicatively coupled. A further hardware-implemented module may then, at a later time, access the memory device to retrieve and process the stored output. Hardware-implemented modules may also initiate communications with input or output devices, and can operate on a resource (e.g., a collection of information).

(177) The various operations of example methods described herein may be performed, at least partially, by one or more processors that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Whether temporarily or permanently configured, such processors may constitute processor-implemented modules that operate to perform one or more operations or functions. The modules referred to herein may, in some example embodiments, comprise processor-implemented modules.

(178) Similarly, the methods described herein may be at least partially processor-implemented. For example, at least some of the operations of a method may be performed by one or more processors or processor-implemented modules. The performance of certain of the operations may be distributed among the one or more processors, not only residing within a single machine, but deployed across a number of machines. In some example embodiments, the processor or processors may be located in a single location (e.g., within a home environment, an office environment, or a server farm), while in other embodiments the processors may be distributed across a number of locations.

(179) The one or more processors may also operate to support performance of the relevant operations in a “cloud computing” environment or as a “software as a service” (SaaS). For example, at least some of the operations may be performed by a group of computers (as examples of machines including processors), these operations being accessible via a network (e.g., the Internet) and via one or more appropriate interfaces (e.g., APIs).

(180) Electronic Apparatus and System

(181) Example embodiments may be implemented in digital electronic circuitry, or in computer hardware, firmware, or software, or in combinations of them. Example embodiments may be implemented using a computer program product, e.g., a computer program tangibly embodied in an information carrier, e.g., in a machine-readable medium for execution by, or to control the operation of, data processing apparatus, e.g., a programmable processor, a computer, or multiple computers.

(182) A computer program can be written in any form of programming language, including compiled or interpreted languages, and it can be deployed in any form, including as a standalone program or as a module, subroutine, or other unit suitable for use in a computing environment. A computer program can be deployed to be executed on one computer or on multiple computers at one site or distributed across multiple sites and interconnected by a communication network.

(183) In example embodiments, operations may be performed by one or more programmable processors executing a computer program to perform functions by operating on input data and generating output. Method operations can also be performed by, and apparatus of example embodiments may be implemented as, special purpose logic circuitry, e.g., an FPGA or an ASIC.

(184) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other. In embodiments deploying a programmable computing system, it will be appreciated that both hardware and software architectures merit consideration. Specifically, it will be appreciated that the choice of whether to implement certain functionality in permanently configured hardware (e.g., an ASIC), in temporarily configured hardware (e.g., a combination of software and a programmable processor), or in a combination of permanently and temporarily configured hardware may be a design choice. Below are set out hardware (e.g., machine) and software architectures that may be deployed, in various example embodiments.

(185) Example Machine Architecture and Machine-Readable Medium

(186) FIG. 25 is a block diagram of a machine in the example form of a computer system 2500 within which instructions 2524 may be executed for causing the machine to perform any one or more of the methodologies discussed herein. In alternative embodiments, the machine operates as a standalone device or may be connected (e.g., networked) to other machines. In a networked deployment, the machine may operate in the capacity of a server or a client machine in server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine may be a personal computer (PC), a tablet PC, a set-top box (STB), a personal digital assistant (PDA), a cellular telephone, a web appliance, a network router, switch, or bridge, or any machine capable of executing instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while only a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

(187) The example computer system 2500 includes a processor 2502 (e.g., a central processing unit (CPU), a graphics processing unit (GPU), or both), a main memory 2504, and a static memory 2506, which communicate with each other via a bus 2508. The computer system 2500 may further include a video display unit 2510 (e.g., a liquid crystal display (LCD) or a cathode ray tube (CRT)). The computer system 2500 also includes an alphanumeric input device 2512 (e.g., a keyboard or a touch-sensitive display screen), a user interface (UI) navigation (or cursor control) device 2514 (e.g., a mouse), a disk drive unit 2516, a signal generation device 2518 (e.g., a speaker), and a network interface device 2520.

(188) Machine-Readable Medium

(189) The disk drive unit 2516 includes a machine-readable medium 2522 on which is stored one or more sets of data structures and instructions 2524 (e.g., software) embodying or utilized by any one or more of the methodologies or functions described herein. The instructions 2524 may also reside, completely or at least partially, within the main memory 2504 and/or within the processor 2502 during execution thereof by the computer system 2500, with the main memory 2504 and the processor 2502 also constituting machine-readable media 2522.

(190) While the machine-readable medium 2522 is shown in an example embodiment to be a single medium, the term “machine-readable medium” may include a single medium or multiple media (e.g., a centralized or distributed database, and/or associated caches and servers) that store the one or more instructions 2524 or data structures. The term “machine-readable medium” shall also be taken to include any tangible medium that is capable of storing, encoding, or carrying instructions 2524 for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure, or that is capable of storing, encoding, or carrying data

structures utilized by or associated with such instructions 2524. The term “machine-readable medium” shall accordingly be taken to include, but not be limited to, solid-state memories, and optical and magnetic media. Specific examples of machine-readable media 2522 include non-volatile memory, including by way of example semiconductor memory devices, e.g., erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), and flash memory devices; magnetic disks such as internal hard disks and removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks.

(191) Transmission Medium

(192) The instructions 2524 may further be transmitted or received over a communications network 2526 using a transmission medium. The instructions 2524 may be transmitted using the network interface device 2520 and any one of a number of well-known transfer protocols (e.g., HTTP). Examples of communication networks include a local area network (LAN), a wide area network (WAN), the Internet, mobile telephone networks, plain old telephone (POTS) networks, and wireless data networks (e.g., WiFi and WiMax networks). The term “transmission medium” shall be taken to include any intangible medium that is capable of storing, encoding, or carrying instructions 2524 for execution by the machine, and includes digital or analog communications signals or other intangible media to facilitate communication of such software.

(193) Although an embodiment has been described with reference to specific example embodiments, it will be evident that various modifications and changes may be made to these embodiments without departing from the broader spirit and scope of the disclosure. Accordingly, the specification and drawings are to be regarded in an illustrative rather than a restrictive sense. The accompanying drawings that form a part hereof show by way of illustration, and not of limitation, specific embodiments in which the subject matter may be practiced. The embodiments illustrated are described in sufficient detail to enable those skilled in the art to practice the teachings disclosed herein. Other embodiments may be utilized and derived therefrom, such that structural and logical substitutions and changes may be made without departing from the scope of this disclosure. This Detailed Description, therefore, is not to be taken in a limiting sense, and the scope of various embodiments is defined only by the appended claims, along with the full range of equivalents to which such claims are entitled.

(194) Such embodiments of the inventive subject matter may be referred to herein, individually and/or collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any single invention or inventive concept if more than one is in fact disclosed. Thus, although specific embodiments have been illustrated and described herein, it should be appreciated that any arrangement calculated to achieve the same purpose may be substituted for the specific embodiments shown. This disclosure is intended to cover any and all adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent to those of skill in the art upon reviewing the above description.

Claims

1. A computing system for utilizing relationship data, the computing system comprising: at least one processor programmed to perform operations comprising: receiving, from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship, and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure at least in part by determining that the first

plurality of confirmation points-in-time do not extend past at least one of the beginning or the end of the first time period; determining a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query to the second computing system.

2. The computing system of claim 1, the operations further comprising: determining that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and determining that a number of the first plurality of confirmation points-in-time are before the first test point-in-time, the response to the first time-dependent query indicating that the first test relationship is true at the first test point-in-time.

3. The computing system of claim 1, the operations further comprising: determining that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and fitting a probability distribution to the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

4. The computing system of claim 1, the operations further comprising: determining that the beginning of the first time period is not defined by the knowledge graph data structure; determining that the end of the first time period is not defined by the knowledge graph data structure; and positioning a probability distribution using the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

5. The computing system of claim 1, the operations further comprising: determining that the first test relationship is a multi-period relationship; selecting a first confirmation point-in-time of the first plurality of confirmation points-in-time that is after the first test point-in-time; and selecting a second confirmation point-in-time of the first plurality of confirmation points-in-time that is before the first test point-in-time, the veracity of the first test relationship at the first test point-in-time being based at least in part on a first probability distribution centered at the first confirmation point-in-time and a second probability distribution centered on the second confirmation point-in-time.

6. The computing system of claim 1, the accessing of the confirmation data comprising accessing a first triple data unit from the knowledge graph data structure, the first triple data unit corresponding to the first test relationship.

7. The computing system of claim 1, the accessing of the confirmation data comprising accessing a plurality of triple data units from the knowledge graph data structure.

8. The computing system of claim 1, the operations further comprising: accessing a candidate relationship comprising a candidate subject entity, a candidate predicate, a candidate object, and a candidate point-in-time; accessing a second plurality of confirmation points-in-time at which the candidate relationship is true; determining veracity of the candidate relationship at the candidate point-in-time using the second plurality of confirmation points-in-time; determining that the veracity of the candidate relationship at the candidate point-in-time is greater than a threshold value; and after determining that the veracity of the candidate relationship at the candidate point-in-time is greater than the threshold value, incorporating the candidate relationship into the knowledge graph data structure.

9. The computing system of claim 1, the operations further comprising: accessing a second test triple data unit from the knowledge graph data structure, the second test triple data unit describing a second test relationship and a second test point-in-time; accessing second confirmation data from the knowledge graph data structure, the second confirmation data describing a second plurality of confirmation points-in-time at which the second test relationship is true; determining that a veracity of the second test relationship at the second test point-in-time is less than a threshold value; and obtaining a confirmation of the second test relationship.

10. The computing system of claim 1, the knowledge graph data structure being stored at a first

tenancy of a plurality of tenancies at a cloud deployment, the operations further comprising: accessing a second knowledge graph data structure stored at a second tenancy of the plurality of tenancies at the cloud deployment; and using the second knowledge graph data structure to determine at least a portion of the first plurality of confirmation points-in-time at which the first test relationship is true.

11. A method for utilizing relationship data in a knowledge graph computing system, the method comprising: receiving, by the knowledge graph computing system and from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing, by the knowledge graph computing system, confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure at least in part by determining that the first plurality of confirmation points-in-time do not extend past at least one of the beginning or the end of the first time period; determining, by the knowledge graph computing system, a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query, by the knowledge graph computing system and to the second computing system.

12. The method of claim 11, further comprising: determining, by the knowledge graph computing system, that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and determining, by the knowledge graph computing system, that a number of the first plurality of confirmation points-in-time are before the first test point-in-time, the response to the first time-dependent query indicating that the first test relationship is true at the first test point-in-time.

13. The method of claim 11, further comprising: determining, by the knowledge graph computing system, that the beginning of the first time period is defined and that the end of the first time period is not defined by the knowledge graph data structure; and fitting a probability distribution to the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

14. The method of claim 11, further comprising: determining, by the knowledge graph computing system, that the beginning of the first time period is not defined by the knowledge graph data structure; determining, by the knowledge graph computing system, that the end of the first time period is not defined by the knowledge graph data structure; and positioning a probability distribution using the first plurality of confirmation points-in-time, the veracity of the first test relationship at the first test point-in-time being a value of the probability distribution at the first test point-in-time.

15. The method of claim 11, further comprising: determining, by the knowledge graph computing system, that the first test relationship is a multi-period relationship; selecting, by the knowledge graph computing system, a first confirmation point-in-time of the first plurality of confirmation points-in-time that is after the first test point-in-time; and selecting, by the knowledge graph computing system, a second confirmation point-in-time of the first plurality of confirmation points-in-time that is before the first test point-in-time, the veracity of the first test relationship at the first test point-in-time being based at least in part on a first probability distribution centered at the first confirmation point-in-time and a second probability distribution centered on the second confirmation point-in-time.

16. The method of claim 11, the accessing of the confirmation data comprising accessing a first triple data unit from the knowledge graph data structure, the first triple data unit corresponding to the first test relationship.

17. The method of claim 11, further comprising: accessing, by the knowledge graph computing system, a candidate relationship comprising a candidate subject entity, a candidate predicate, a candidate object, and a candidate point-in-time; accessing, by the knowledge graph computing system, a second plurality of confirmation points-in-time at which the candidate relationship is true; determining, by the knowledge graph computing system, a veracity of the candidate relationship at the candidate point-in-time using the second plurality of confirmation points-in-time; determining, by the knowledge graph computing system, that the veracity of the candidate relationship at the candidate point-in-time is greater than a threshold value; and after determining that the veracity of the candidate relationship at the candidate point-in-time is greater than the threshold value, incorporating the candidate relationship into the knowledge graph data structure.

18. The method of claim 11, further comprising: accessing, by the knowledge graph computing system, a second test triple data unit from the knowledge graph data structure, the second test triple data unit describing a second test relationship and a second test point-in-time; accessing, by the knowledge graph computing system, second confirmation data from the knowledge graph data structure, the second confirmation data describing a second plurality of confirmation points-in-time at which the second test relationship is true; determining, by the knowledge graph computing system, that a veracity of the second test relationship at the second test point-in-time is less than a threshold value; and obtaining, by the knowledge graph computing system, a confirmation of the second test relationship.

19. The method of claim 11, the knowledge graph data structure being stored at a first tenancy of a plurality of tenancies at a cloud deployment, the method further comprising: accessing, by the knowledge graph computing system, a second knowledge graph data structure stored at a second tenancy of the plurality of tenancies at the cloud deployment; and using, by the knowledge graph computing system, the second knowledge graph data structure to determine at least a portion of the first plurality of confirmation points-in-time at which the first test relationship is true.

20. A machine-readable medium comprising instructions thereon that, when executed by at least one processor, cause the at least one processor to perform operations comprising: receiving, from a second computing system, a first time-dependent query against a knowledge graph data structure, the first time-dependent query comprising an indication of a first test relationship, and an indication of a first test point-in-time, the first test relationship being described by a first time period during which the first test relationship is true; accessing confirmation data from the knowledge graph data structure, the confirmation data describing a first plurality of confirmation points-in-time at which the first test relationship is true; determining that at least one of a beginning or an end of the first time period is not defined by the knowledge graph data structure at least in part by determining that the first plurality of confirmation points-in-time do not extend past at least one of the beginning or the end of the first time period; determining a response to the first time-dependent query indicating a veracity of the first test relationship at the first test point-in-time using the first plurality of confirmation points-in-time; and sending the response to the first time-dependent query to the second computing system.
