

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12406060
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Gupta; Satya V.

Automated interpreted application control for workloads

Abstract

Embodiments provide functionality to protect computing workloads from script-based attacks. Upon receipt, at a workload, of a command to commence execution of code of a script, an embodiment determines whether (i) permissions of a user issuing the command comply with a permissions security standard, (ii) an identifier of an interpreter supporting the script is included in an approved interpreter list, (iii) an identifier of a selected parameter of the interpreter is included in an approved parameter list, and (iv) an identifier of the script is included in an approved list of executables. If all of the aforementioned checks pass, such an embodiment allows execution of the code of the script; otherwise, execution is denied, thereby protecting the workload in an event of a script-based attack.

Inventors:	Gupta; Satya V. (Dublin, CA)
Applicant:	Virsec Systems, Inc. (San Jose, CA)
Family ID:	82021126
Assignee:	Virsec Systems, Inc. (San Jose, CA)
Appl. No.:	17/664009
Filed:	May 18, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20220391506 A1	Dec. 08, 2022

Related U.S. Application Data

us-provisional-application US 63190100 20210518
us-provisional-application US 63190099 20210518

Publication Classification

Int. Cl.: G06F21/56 (20130101); G06F21/57 (20130101)

U.S. Cl.:

CPC G06F21/563 (20130101); G06F21/562 (20130101); G06F21/566 (20130101); G06F21/577 (20130101); G06F2221/033 (20130101)

Field of Classification Search

CPC: G06F (21/562); G06F (21/566); G06F (21/577); G06F (2221/033); G06F (21/563)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6321334	12/2000	Jerger	726/1	G06F 21/54
8347267	12/2012	Givoni et al.	N/A	N/A
9246935	12/2015	Lietz et al.	N/A	N/A
9374390	12/2015	Teal	N/A	G06F 21/554
9418230	12/2015	Archer et al.	N/A	N/A
9501650	12/2015	Chess et al.	N/A	N/A
9578060	12/2016	Brisebois et al.	N/A	N/A
9589560	12/2016	Vitaladevuni et al.	N/A	N/A
10116681	12/2017	Cornell et al.	N/A	N/A
10354074	12/2018	Gupta	N/A	N/A
10387659	12/2018	Youngberg et al.	N/A	N/A
10395041	12/2018	Youngberg et al.	N/A	N/A
10447730	12/2018	Gupta	N/A	N/A
10467419	12/2018	Youngberg et al.	N/A	N/A
10963565	12/2020	Xu et al.	N/A	N/A
11615061	12/2022	Malik et al.	N/A	N/A
11907378	12/2023	Gupta	N/A	N/A
12158958	12/2023	Gupta et al.	N/A	N/A
12282552	12/2024	Gupta	N/A	N/A
2003/0074207	12/2002	Pace et al.	N/A	N/A
2004/0046785	12/2003	Keller	N/A	N/A
2005/0102534	12/2004	Wong	726/4	G06F 21/577
2007/0199000	12/2006	Shekhel	719/330	G06F 9/468
2008/0052527	12/2007	Siedlarz	N/A	N/A
2009/0119769	12/2008	Ross et al.	N/A	N/A
2013/0111595	12/2012	Amit et al.	N/A	N/A
2014/0082735	12/2013	Beskrovny et al.	N/A	N/A
2014/0082739	12/2013	Chess et al.	N/A	N/A
2014/0165192	12/2013	Zhu et al.	N/A	N/A
2015/0215332	12/2014	Curcic et al.	N/A	N/A
2015/0261653	12/2014	Lachambre et al.	N/A	N/A
2015/0309813	12/2014	Patel	N/A	N/A

2016/0164891	12/2015	Satish et al.	N/A	N/A
2016/0241582	12/2015	Boia et al.	N/A	N/A
2017/0270303	12/2016	Roichman et al.	N/A	N/A
2017/0288878	12/2016	Lee et al.	N/A	N/A
2017/0353434	12/2016	Al-Saber et al.	N/A	N/A
2018/0349602	12/2017	Johns	N/A	N/A
2019/0138725	12/2018	Gupta	N/A	N/A
2019/0286833	12/2018	Takumi et al.	N/A	N/A
2019/0377877	12/2018	Johns	N/A	N/A
2020/0004963	12/2019	Zheng et al.	N/A	N/A
2020/0042714	12/2019	Gupta	N/A	N/A
2020/0065166	12/2019	Myneni et al.	N/A	N/A
2020/0097662	12/2019	Hufsmith et al.	N/A	N/A
2020/0134193	12/2019	Youngberg et al.	N/A	N/A
2020/0134194	12/2019	Youngberg et al.	N/A	N/A
2020/0134195	12/2019	Youngberg et al.	N/A	N/A
2020/0167477	12/2019	Ionescu et al.	N/A	N/A
2020/0336507	12/2019	Lee et al.	N/A	N/A
2021/0086089	12/2020	Pardeshi et al.	N/A	N/A
2021/0099483	12/2020	Shukla	N/A	N/A
2021/0160273	12/2020	Choi et al.	N/A	N/A
2022/0046031	12/2021	Kaidi	N/A	N/A
2022/0067174	12/2021	Gupta	N/A	N/A
2022/0147635	12/2021	Copty et al.	N/A	N/A
2022/0198025	12/2021	Gupta et al.	N/A	N/A
2022/0207151	12/2021	Gupta	N/A	N/A
2022/0210180	12/2021	Gupta	N/A	N/A
2022/0214928	12/2021	Gupta et al.	N/A	N/A
2023/0004652	12/2022	Gupta	N/A	N/A
2025/0068726	12/2024	Gupta et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2015/007166	12/2014	WO	N/A
2016/130372	12/2015	WO	N/A
2020/026228	12/2019	WO	N/A
2022/047245	12/2021	WO	N/A
2022/147474	12/2021	WO	N/A
2022/147478	12/2021	WO	N/A
2022/155685	12/2021	WO	N/A
2022/155687	12/2021	WO	N/A
2022/246436	12/2021	WO	N/A
2022/246437	12/2021	WO	N/A
2023/133586	12/2022	WO	N/A

OTHER PUBLICATIONS

Dizdar, A. , “Why are SAST solutions not always the best option for AST?”, Bright, Jan. 22, 2020, 10 pages. cited by applicant

Horvath, M., et al., “Magic Quadrant for Application Security Testing” Gartner Reprint, (33 pages) Apr. 29, 2020, 33 pages, retrieved from <https://www.gartner.com/doc/reprints?id=1->

1YWZKUB5&ct=200429&st=sb. cited by applicant

Potdar, S., "The Curious Case of False Positives in Application Security", Security Zone, May 13, 2019, 4 pages. cited by applicant

QwietAI "ShiftLeft Achieves Highest Ever SAST Score on OWASP Benchmark", Nov. 1, 2023, 7 pages, retrieved from <https://qwiet.ai/news-press/shiftright-achieves-highest-ever-sast-score-on-owasp-benchmark/>. cited by applicant

"ATT&CK (Registered)", The Mitre, Available online at: <<https://attack.mitre.org>>, Sep. 14, 2023, 1 page. cited by applicant

"CAPEC—Common Attack Pattern Enumeration and Classification", (CAPEC(Trademark)), Available on <https://capec.mitre.org>, Jan. 2, 2022, 2 pages. cited by applicant

"CAPEC View: Mechanisms of Attack", View ID: 1000, (Version 3.6), Available online <https://capec.mitre.org/data/definitions/1000.html>, Jan. 2, 2022, 2 pages. cited by applicant

"Configuration Management Database (CMDB)", 2023, 15 pages. cited by applicant

"CWE View: Software Development" View ID: 699, Individual Dictionary Definition (4.6), Available on <https://cwe.mitre.org/data/definitions/699.html>, Jan. 2, 2022, 2 pages. cited by applicant

"Integrated Risk Management (IRM)", Gartner Glossary, Available on <https://www.gartner.com/en/information-technology/glossary/integrated-risk-management-irm>, Jan. 2, 2022, 5 pages. cited by applicant

"javap—The Java Class File Disassembler", Oracle, Java SE Documentation, Available on <https://docs.oracle.com/javase/7/docs/technotes/tools/windows/javap.html>, Jan. 2, 2022, 3 pages. cited by applicant

"The Shadow Brokers", Wikipedia, Available on https://en.wikipedia.org/wiki/The_Shadow_Brokers, Jan. 2, 2022, 8 pages. cited by applicant

"Welcome to YARA's documentation-yara 4.3.2 documentation" Available online at: <<https://yara.readthedocs.io/en/stable/>>, retrieved on Sep. 14, 2023, 4 pages. cited by applicant

"Windows Print Spooler Remote Code Execution Vulnerability", Security Vulnerability, CVE-2021-34527, Available on <https://msrc.microsoft.com/update-guide/vulnerability/CVE-2021-34527#title>, Jun. 13, 2023, 1 page. cited by applicant

Caputo, D., et al: "Droids in Disarray: Detecting Frame Confusion in Hybrid Android Apps", Jun. 11, 2019 (Jun. 11, 2019), Advances in Databases and Information Systems; [Lecture Notes in Computer Science; Lect.Notes Computer], Springer International Publishing, Cham, pp. 121-139. cited by applicant

Chen et al., "Automated system change discovery and management in the cloud", IBM Journal of Research and Development, vol. 60, No. 2-3, Mar. 1, 2016, pp. 2:1-2:10. cited by applicant

Chen et al., "Detecting and Identifying System Changes in the Cloud via Discovery by Example", 2014 IEEE International Conference on Big Data, Oct. 27, 2014, pp. 90-99. cited by applicant

Chen, Lu., et al., "Research on Mobile Application Local Denial of Service Vulnerability Detection Technology Basec on Rule Matching", 2019 IEEE International Conference On Energy Internet (ICEI), IEEE, May 27, 2019 (May 27, 2019), pp. 585-590. cited by applicant

David Stahl, "What's an RFC and what can they do for me?", Global Knowledge, Available on <https://www.globalknowledge.com/us-en/resources/resource-library/articles/whats-an-rfc-and-what-can-they-do-for-me/>, Sep. 16, 2009, 2 pages. cited by applicant

International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2021/048077, mailed on Dec. 20, 2021, 11 pages. cited by applicant

International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2021/073197, mailed on Mar. 21, 2022, 10 pages. cited by applicant

International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2021/073201, mailed on Apr. 7, 2022, 10 pages. cited by applicant

International Search Report and Written Opinion received for PCT Patent Application No.

PCT/US2022/070236, mailed on Apr. 8, 2022, 15 pages. cited by applicant
International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2022/070240, mailed on Apr. 7, 2022, 17 pages. cited by applicant
International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2022/072416, mailed on Sep. 8, 2022, 13 pages. cited by applicant
International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2022/072417, mailed on Sep. 8, 2022, 15 pages. cited by applicant
Johns, M., et al., "XSSDS: Server-Side Detection of Cross-Site Scripting Attacks", Computer Security Applications Conference, 2008. Acsac, Dec. 8, 2008, pp. 335-344. cited by applicant
Koutroumpouchos, K., et al., "ObjectMap: detecting insecure object deserialization", PCI '19: Proceedings of the 23rd Pan-Hellenic Conference on Informatics, Nov. 28, 2019, pp. 67-72. cited by applicant
Prevelakis et al., "Sandboxing Applications", USENIX, The Advanced Computing Systems Association, Feb. 25, 2019, pp. 1-9. cited by applicant
Secure execution of privileged scripts ED-Darl Kuhn, IP.Com, IP.Com Inc., West Henrietta, Sep. 18, 2009, XP013134389. cited by applicant
Sun, F., et al., "Client-Side Detection of XSS Worms by Monitoring Payload Propagation", Advances In Databases And Information Systems, Sep. 21, 2009, pp. 539-554. cited by applicant
Xia, M. et al: "Effective Real-Time Android Application Auditing", 2015 IEEE Symposium On Security and Privacy, IEEE, May 17, 2015 (May 17, 2015), pp. 899-914. cited by applicant
International Search Report and Written Opinion received for PCT Patent Application No. PCT/US2023/060379, mailed on Apr. 19, 2023, 15 pages. cited by applicant
Yang Ji, "Efficient and refinable attack investigation," A Dissertation Presented to The Academic Faculty, Georgia Institute of Technology, Dec. 2019, pp. 123. cited by applicant

Primary Examiner: Callahan; Paul E

Attorney, Agent or Firm: Hamilton, Brook, Smith & Reynolds, P.C.

Background/Summary

RELATED APPLICATIONS (1) This application claims the benefit of U.S. Provisional Application No. 63/190,099, filed on May 18, 2021, and U.S. Provisional Application No. 63/190,100, filed on May 18, 2021. The entire teachings of the above applications are incorporated herein by reference.

BACKGROUND

(1) Workloads utilize various computing resources to load and execute appropriate software instructions so as to accomplish user desired tasks. Workloads may be deployed across a network of an organization such as an enterprise, and may feature, for example, various versions of sets of software instructions.

(2) Cyber-attacks inflicted upon such workloads are becoming increasingly sophisticated. Attacks are often targeted to exploit specific vulnerabilities. Various methods and tools exist for identifying these vulnerabilities, but these existing methods and tools are inadequate.

SUMMARY

(3) Embodiments provide functionality to protect computing workloads from script-based attacks.

(4) This disclosure provides a method where a workload receives a command to commence execution of code of a script. The method determines whether (i) permissions of a user issuing the command comply with a permissions security standard, (ii) an identifier of an interpreter

supporting the script is included in an approved interpreter list, (iii) an identifier of a selected parameter of the interpreter is included in an approved parameter list, and (iv) an identifier of the script is included in an approved list of executables. The method protects the workload from script-based attacks by allowing or denying the execution of the code of the script based on the determining.

(5) A permissions security standard, for example, in Linux, may specify that the user issuing the command must have ownership permission of the script, or at least membership of a user group associated with the script, in order for execution of code of the script to be allowed. Such a security standard may prevent a user from relying on the permission of “others,” i.e., non-owners and non-group-members associated with the script, in order to run code of the script. Embodiments may utilize any desired permissions security standard. An approved interpreter list may include multiple levels of granularity, such as by considering specific versions of noted interpreters. The selected parameter under consideration may be included in the command, and may serve to modify operations performed by the interpreter as provided by a developer of the interpreter and as desired by the user. Such parameters may alternatively be referred to as command line switches, or simply as switches. At least one of the identifier of the interpreter, the identifier of the selected parameter, and the identifier of the script may include a checksum or another measure/indication of script integrity. In an example embodiment, all such interpreter switches/parameters are checked to determine if the selected parameters comply with user/enterprise requirements.

(6) The method may include receiving an identifier of an approved interpreter, or parameter thereof, from a vendor of the approved interpreter, and adding the identifier of the approved interpreter, or parameter thereof to the approved interpreter list, or the approved parameter list. Alternatively, or in addition, the method may include receiving an identifier of an approved script from a vendor of the approved script, and adding the identifier of the approved script to the approved list of executables.

(7) Allowing or denying the execution of the code of the script may include denying the execution of the code of the script in response to the received command not being present in a list of commands known to be used by the workload at runtime. Allowing or denying the execution of the code of the script may be further based on determining whether the user is positively authorized to execute the script based on permissions of the user and permissions of the script. Further still, allowing or denying the execution of the code of the script may include denying the execution of the code of the script in response to a process to be spawned by the script matching a process in a list of known malicious processes. Allowing or denying the execution of the code of the script may further include denying the execution of the code of the script in response to at least one of the script, a process to be spawned by the script, and a library to be loaded during the execution of the script, including a known vulnerability.

(8) In some implementations, in response to allowing the execution of the code of the script, the method intercepts a given process spawned during the execution of the code of the script, tokenizes the received command, compares tokens of the tokenized received command with entries in an allow-list database, and, based on the comparing, maintains the execution, or suspends the execution, of the code of the script, or indicates a result of the comparing to a user seeking to control execution of the script.

(9) In response to allowing the execution of the code of the script, the method may include intercepting a library loaded during the execution of the script. In such an implementation, in response to the library including a known vulnerability, and a disposition of the library regarding the vulnerability having not been received from a user, the method further includes suspending the execution of the code of the script, or indicates the vulnerability to a user seeking to control execution of the script.

(10) Another aspect of the disclosure is directed to a system for protecting a computing workload from script-based attacks. The system includes an application control policy (ACP) engine and an

endpoint process monitoring client (EPMC). The EPMC may be configured to intercept a command to commence execution of code of a script. The ACP engine may be configured to determine, in response to the interception of the command by the EPMC, whether (i) permissions of a user issuing the command comply with a permissions security standard, (ii) an identifier of an interpreter supporting the script is included in an approved interpreter list, (iii) an identifier of a selected parameter of the interpreter is included in an approved parameter list, and (iv) an identifier of the script is included in an approved list of executables. The ACP engine may be further configured to allow or deny the execution of the code of the script based on the determining, thereby protecting the workload from script-based attacks.

(11) Yet another aspect of the disclosure is directed to a computer program product for protecting a computing workload from script-based attacks. The computer program product includes one or more non-transitory computer-readable storage devices and program instructions stored on at least one of the one or more storage devices. The program instructions, when loaded and executed by a processor, cause the processor to protect a computing workload from script-based attacks as described herein.

(12) It is noted that embodiments of the method, system, and computer program product may be configured to implement any embodiments described herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The foregoing will be apparent from the following more particular description of example embodiments, as illustrated in the accompanying drawings in which like reference characters refer to the same parts throughout the different views. The drawings are not necessarily to scale, emphasis instead being placed upon illustrating embodiments.

(2) FIG. 1 is a block diagram illustrating a legacy or monolith service web application that may be analyzed using embodiments.

(3) FIG. 2 is a block diagram illustrating a cloud native or microservice web application that may be analyzed using embodiments.

(4) FIG. 3 is a schematic block diagram representation of an example software infrastructure subject to embodiments.

(5) FIG. 4 is a block diagram showing an example workload subject to embodiments.

(6) FIG. 5 is a flowchart of a method of protecting a computing workload from script-based attacks according to an embodiment.

(7) FIG. 6 is a simplified block diagram of a system for protecting a computing workload from script-based attacks according to an embodiment.

(8) FIG. 7 illustrates a computer network or similar digital processing environment in which embodiments may be implemented.

(9) FIG. 8 is a diagram illustrating an example internal structure of a computer in the environment of FIG. 7.

DETAILED DESCRIPTION

(10) A description of example embodiments follows.

(11) Cyber-attackers have started graduating from unleashing file-based, compiled malware to relying on file-less, script-based malware. This newer type of malware executes within context of an executable, such as an interpreter, that is already present on a computing workload. This type of script-based attack may also be referred to as a “living off the land” type of attack, due to a dual-use nature of interpreters provided as part of an operating system installation or distribution. Such a dual-use nature permits these interpreters to be used by information technology personnel of an enterprise for legitimate activities, but also leaves such interpreters vulnerable to use by attackers

for nefarious purposes. As interpreters thus used illicitly are otherwise considered legal in and of themselves, it is important to examine the manner in which such interpreters are used, as described herein.

(12) Many existing file reputation engines classify malware based on a signature such as a checksum of an executable. This checksum is relatively immutable. Cybersecurity providers promoting these types of engines claim they have acquired signatures of billions of malware manifestations through research from numerous anti-virus vendors. When an end-user of such a file reputation engine uploads either a given piece of malware or its checksum, the file reputation engine responds with the name of the piece of malware and a score indicating how many different antivirus engines correctly identify that piece of malware.

(13) File reputation engines such as these use a rules engine, e.g., a YARA rules engine (available at <https://yara.readthedocs.io/en/stable/>), which relies on a database of strings that are embedded in the executable associated with the named piece of malware. These strings can be of the following types: a) Hexadecimal strings b) Text strings (e.g., case insensitive strings, wide character strings, XOR strings, base-64 strings, full word strings, etc.) c) Regular expressions d) Private strings

(14) Unfortunately, existing file reputation engines are nearly helpless against file-less, or script-based malware. Many bad actors take advantage of this fact to first exploit a vulnerability in a workload and then establish a connection, such as a reverse shell, to their command control server. Attackers can then send more script-based malware through the established connection, while evading detection on a persistent basis.

(15) The foregoing factors undermine the risk posture of an enterprise. Embodiments described herein address the abovementioned shortcomings and allow users to block bad actors from exploiting vulnerabilities and launching script-based attacks using web shells or similar tactics. Many such tactics, from which embodiments offer protection, are disclosed in MITRE's ATT&CK framework (available at <https://attack.mitre.org/>).

(16) Legacy And Cloud Native Applications

(17) Enterprises have been slowly migrating from legacy or monolith applications to cloud native or microservices-based applications. As the name suggests, instead of deploying a large monolith application, users of microservices-based applications can disaggregate the overall functionality of a monolith application into smaller, but modular chunks. This disaggregation allows microservices-based applications to be upgraded and patched without deploying the full application all over again. Another advantage of microservices-based applications is that a hot microservice can scale independently of the other microservices that are not invoked as much. Scaling the entire monolith application is also not desirable.

(18) A legacy, i.e., monolith, application **100** is represented in FIG. 1. In this monolith application **100** represented in FIG. 1, all individual functionality (operations **102a**, access **102b**, node manager **102c**, messenger **102d**, distributed tracing **102e**, and topology **102f**) is packaged in one monolith workload **103**. This common functionality **102a-f** shares a local or remote common database **104** workload as well. As such, each individual function **102a-d** sends respective data **105a**, **105b**, **105ca**, **105cb**, and **105d** to the common database **104**. Intercommunication between discrete functionality **102a-d** in the monolith application **100** is typically performed via a very fast inter-process communication (IPC) mechanism such as shared memory or pipes/fifos. This is possible since the entire functionality **102a-f** of the workload **103** lives on one host **106**.

(19) In operation, the application **100** is accessed by the client applications which include a mobile application **101a**, traditional web application **101b**, and single-page web application **101c**. The mobile application **101a** and single-page web application **101c** access the functionality **102a-f** of the application **100** via the application programming interface (API) gateway **107**. The traditional web application **101b** accesses the application **100** via the web application service **108** and API gateway **107**.

(20) The monolith application **100** runs on top of operating system (OS) provided runtime

processes and services. In addition, the application **100** may leverage third party binary and interpreted code (executables and libraries) as well as first party code (executables and libraries). All this code, i.e., third party binary code, third party interpreted code, and first party code, may contain vulnerabilities, some of which are known, and others that may be latent or unknown at a given moment in time.

(21) The application **100** can be delivered as a microservices-based application. FIG. **2** illustrates an example application **220**, which is a microservices-based version of the application **100**. As can be seen, the functionality **102a-f** of the monolith application **100** is broken into individual services **222a-f** (i.e., containers) with individual databases **224a-d** and web connections **229a-b**. Such web connections **229a-b** may permit a service **222e-f** to forward a request to a remote database or to another service via the Internet or other such network connection. Each functionality **223a-f** can live on one or more host **226** or workload **223a-f**. The communication between individual microservices **222a-f** is via messages **225a-f** that are brokered by an event bus **230**. Individual user transactions weave through one or more microservices **222a-f**, and when the full transaction is completed successfully, each service involved in the transaction synchronizes its respective database **224a-d** atomically. Conversely, if the transaction fails, changes to the individual databases **224a-d** are not committed.

(22) Similar to the application **100**, in operation the application **220** is accessed by the client applications, which include a mobile application **101a**, traditional web application **101b**, and single-page web application **101c**. The mobile application **101a** and single-page web application **101c** access the functionality **222a-f** of the application **220** via the application programming interface (API) gateway **107**. The traditional web application **101b** accesses the application **220** via the web application service **108** and API gateway **107**.

(23) How to Evaluate Cyber Risk of an Application

(24) Software infrastructure, e.g., the software infrastructure utilized by an organization, can be visualized as being segmented at different tiers such as on a per business unit basis, per application within that business unit, per private/public cloud provider or data center. The application's infrastructure may further be segmented based on a location within the data center and a subnet in said data center. Workloads themselves can be individual workloads or a high-availability or load balanced cluster of identical workloads.

(25) An example of such a visualization is shown in FIG. **3**. Particularly, FIG. **3** illustrates software application infrastructure environment **330** of an enterprise in which embodiments for protecting a computing workload from script-based attacks, as described herein, may be employed.

(26) In the environment **330**, a workload may include a monolith (e.g., **100**) or microservices-based (e.g. **220**) software application. Such an application may be installed at additional workloads deployed across a network.

(27) In FIG. **3**, applications **331a** and **331n**, locations **332a** and **332n**, de-militarized zone (DMZ) **334**, subnet zone **335a** and subnet zone **335n**, services **336a-c**, load balancer **337**, and workloads **338a**, **338b**, **338c**, and **338n** represent network topology of an aspect of the workloads **338a**, **338b**, **338c**, and **338n**, such as an application. Depicted in the lowest layer of the network topology are individual workloads **338a**, **338b**, **338c**, and **338n** that provide the application functionality. Such individual workloads **338a**, **338b**, **338c**, and **338n** may comprise three layers, including an infrastructure layer, a virtualization layer, and a service layer as described hereinbelow in relation to FIG. **4**. In such an embodiment, code used within a given workload **338a**, **338b**, **338c**, and **338n** can be resident in either the file system or in memory. The network environment **330** may include an intranet **339** connected to the Internet **340** and as such may be accessed by an end-user **341**. In some cases, the end-user **341** may be a malicious attacker.

(28) Continuing with respect to FIG. **3**, deployed upon the intranet **339** is business logic for respective business units, which may include a first business unit **342a** and other business units up to and including a Tth business unit **342n**. Such business units **342a** and **342n** may also be referred

to as tenants. Within the business logic for the business units **342a**, **342n** are software applications **331a**, **331n**. While only a first application **331a** and second application **331n** are depicted, the business units **342a** and **342n** may utilize any number of applications. Each such application **331a**, **331n** is deployed on at least one location **332a**, **332n**, which may be a cloud location, on premises location, or other such location known in the art. Within the locations **331a**, **331n** is deployed a DMZ **334**, beyond which are deployed at least one subnet from a first subnet zone **335a** to a Zth subnet zone **335n**. Various services **336a-c** are deployed within the subnets **335a**, **335n**. In particular, services **336a** and **336b** are deployed on subnet zone **335a** and service **336c** is deployed on subnet zone **335n**. Within each service **336a-c**, are deployed workloads **338a-n**. The service **336a** is implemented using the workload **338a** and the service **336b** is implemented using the workload **338b**. Likewise, the service **336c** is deployed using both the workloads **338c** and **338n** and the load balancer **337**. Upon each workload **338a-n**, one or more application service instance (not shown) may be deployed. Each application service instance includes an infrastructure hardware layer, a virtualization layer, and a service, which may include operating system runtime packages, compatible precompiled binary packages, and compatible byte code packages.

(29) Individual workloads, e.g., the workloads **338a**, **338b**, **338c**, and **338n** are composed of three layers as shown by the example workload **440** in FIG. 4. FIG. 4 illustrates an individual workload **440** that may be deployed upon a network to enable functionality of software such as an application. Such a workload **440** includes an infrastructure layer **441**, a virtualization layer **448**, and a service layer **459**. So configured, such a workload **440** may be referred to as an application service instance (ASI). The infrastructure layer **441** defines attributes such as compute, storage, and host OS attributes. This infrastructure layer **441** can be provided and managed by either a 1.sup.st or 3.sup.rd party cloud provider or a private data center provider. The ASI shown in FIG. 4 includes a collection of components comprising a monolith service or a microservice **470**. Such a collection **470** includes virtual machines **449a**, containers **449b**, and serverless functions **449c**. The ASI shown in FIG. 4 encompasses a workload **440** deployed on a server.

(30) The lowest layer of the workload **440** is called the compute or infrastructure layer **441**. The infrastructure layer **441** includes physical hardware **442**, persistent storage **443** available on the network, a host device **444** with a processor and memory, a physical network interface card **445**, local storage **446**, and a host operating system **447**. As such, the layer **441** is composed of not only the compute **442**, memory **444**, networking **445**, and storage (local **446** and remote **443**) components, but also the host operating system **447**.

(31) The layer in the center of the workload **440** is called the virtualization layer **448**. This layer **448** allows users to specify the virtualization technology that isolates the application's business logic from the underlying infrastructure layer **441**. There are three virtualization options—the virtual machine **449a**, the container **449b**, and the serverless option **449c**.

(32) The virtualization layer **448** may include a hypervisor **450** and a guest entity **451** that may include a virtual processor and memory. The virtual layer **448** may also include a virtual network interface card **452**, a virtual disk **453**, and may have an operating system **454** installed thereupon. The virtual layer **448** also includes, for container applications **449b**, container mounts **455**, container runtime components **456**, and network plugin **457**. The virtualization layer **448** may also include a serverless **449c** function handler **458**.

(33) The workload's virtualization layer **448** defines attributes such as a virtualization type, which may be implemented as a bare metal instance, a virtual machine instance **449a**, a container instance **449b**, or a serverless function **449c**. This layer **448** can be provided and managed by either the 1.sup.st party (where the application and infrastructure are owned and operated by the same entity) or by 3.sup.rd parties (where the application and infrastructure are owned and operated by different entities).

(34) The top layer of the workload **440** is the service layer **459**. In operation, the hypervisor **450** of the virtual layer **448** may, through the operating system **454**, connect to one or more virtual

machines **449a** that are part of the service layer **459**. Such virtual machines **449a** may include handlers **460a**, **460b**, **460c**, **460d**, application programming interface (API) or web logic or databases **461a**, **461b**, third-party binaries **462a**, operating system runtime binaries **463**, web frameworks **464a**, **464b**, binary framework **465a**, operating system services **466**, and process name spaces **467a**, **467b**, **467c**.

(35) In embodiments operating upon software configured as containers **449b**, the service layer **459** includes handlers **460e**, **460f**, API or web logic or database **461c**, web frameworks **464c**, process namespace **467d**, **467e**, third-party binaries **462b**, and binary frameworks **465b**.

(36) In serverless configurations **449c**, a serverless function handler **458** interfaces with handles **460g**, **460h**, respectively through APIs or web or business logic functions **468**, and binary functions **469**.

(37) The service layer **459** describes the logical part of the application, implemented as the monolith or microservice collection **470**, upon the workload **440**. As can be seen in FIG. 4, the memory subsystem is indicated by a collection of processes **467a-e** that execute code **461a-c**, **462a-b**, **463**, **468**, and **469** that may be from one or more packages or frameworks **464a-c** **465a-b**, **466**. These frameworks could be web (or API Frameworks) **464a-c**, Binary **465a-b**, or OS Runtime services (aka Host Framework) **466**. The processes **467a-e** from the one or more frameworks **464a-c** **465a-b**, **466** described above not only load code, but also open handles **460a-h** on non-code files such as configuration files, log files, content files etc. Processes **467a-e** can also have handles open on other workloads that form part of the overall end user application.

(38) The service layer **459** contains active code that provides the application's observable functionality. The service layer **459** can be powered by a mixture of OS and OS-provided runtime services (e.g., a host framework), one or more 1.sup.st or 3.sup.rd party precompiled executables and libraries (e.g., binary frameworks), and one or more 1.sup.st or 3.sup.rd party interpreted code files (e.g., interpreted frameworks).

(39) Automatic and On-Demand Evaluation

(40) In current application implementations, such as those described hereinabove in relation to FIG. 3 and FIG. 4, evaluating cyber risk entails determining the risk from individual workloads. For instance, quantifying the cyber risk for the application **331a** entails aggregating the cyber risk from individual workloads **338a**, **338b**, **338c**, and **338n**. Vulnerability databases leverage a common metric called the Common Vulnerability Security Score (CVSS). A CVSS score between 9.0 and 10.0 is considered Critical. A score between 7.0 to 8.9 is considered High. A Score between 4.0 and 6.9 is considered Medium. Attacks that are conducted from remote networks are considered more lethal than attacks that are conducted from local networks because it is typically more difficult to conduct the attack from inside the enterprise for fear of being caught. Similarly, attacks that require privileges are considered harder to mount since it involves stealing a specific user's credentials.

(41) Automatic Interpreted Application Control Policy (IACP) Generation

(42) Embodiments provide a solution that protects against script-based attacks by automatically generating and imposing a series of IACP rules. FIG. 5 is a flow diagram of a method **550** for protecting workloads against script-based attacks according to an embodiment.

(43) The method **550** begins at step **551** with receiving, at a workload, a command to commence execution of code of a script. Such a workload may be of workloads **338a**, **338b**, **338c**, and **338n** or workload **440**. The method **550** continues by determining whether permissions **552** of a user issuing the command comply with a permissions security standard. If the user permissions do not comply **552n**, the method **550** terminates at step **556** by denying the execution of the code of the script based on the determining **552**. If the user permissions comply with the security standard **552y**, the method **550** continues by determining whether an identifier of an interpreter **553** supporting the script is included in an approved interpreter list. Such an identifier may be, for example, a hash or a checksum of the interpreter. If a match is not found **553n**, the method **550** terminates at step **556** by denying the execution of the code of the script based on the determining

553. If a match is found 553y, the method 550 continues by determining whether an identifier of a selected parameter of the interpreter 554 is included in an approved parameter list. Such an identifier may be, for example, a hash or a checksum of the parameter. If a match is not found 554n, the method 550 terminates at step 556 by denying the execution of the code of the script based on the determining 554. If a match is found 554y, the method 550 continues by determining whether an identifier of the script 555 is included in an approved list of executables. If a match is not found 555n, the method 550 terminates at step 556 by denying the execution of the code of the script based on the determining 555. If a match is found 555y, the method 550 terminates at step 557 by allowing the execution of the code of the script based on the determining 552, 553, 554, 555.

(44) It is noted that while FIG. 5 illustrates step 552 being followed by step 553, step 553 being followed by step 554, and step 554 being followed by step 555, embodiments are not so limited. In particular, the determinations 552-555 may be made in any order and contemporaneously.

(45) The following sections describe several sets of IACP policies and rules that may be employed by embodiments to protect a workload from being attacked. When a new script-based process is about to be spawned, the command line for said process is compared against a series of IACP policies and sub-rules thereof. The underlying IACP rules can result in the process being denied further execution, or in the process being permitted to complete execution. Those IACP Policy and Rules that cause process spawning to be denied are described first, followed by the rules that allow the process spawn to continue.

(46) File Integrity Match Allow Policy and Rules

(47) Developers often add scripts into released compiled software packages, such as RPM Package Manager (RPM) packages, Windows Installer (MSI) packages, self-extracting archive (SEA) packages, TAR packages, etc. Some developers also release source code packages that a user can configure for their use case before the objects are compiled on the user's machine. The process of compilation may result in some additional scripts that can be run on the user's machine. These scripts may be considered "authorized" by embodiments unless specifically blocked by the user. An embodiment extracts all scripts, statically and dynamically, that are produced after source code compilation. Inherently, this policy is an allow policy. In an example embodiment, an identifier of such an authorized script may be determined 554 to be in an approved list of scripts according to the method 550. Alternatively, an identifier of an executable containing such an authorized script may be determined 553 to be in an approved list of executables according to the method 550.

(48) Hardened Deny Policy and Rules

(49) Even if code from a package is present on a workload, it may not be needed (i.e., loaded and executed) at run time. Static and dynamic scripts from packages that are not used at runtime can be hardened out by embodiments. Inherently, this policy is a deny policy.

(50) User Mismatch Deny Policy and Rules

(51) Since scripts require an interpreter to be running, it may be difficult for a developer to determine what access privileges should be assigned to the scripts the developer packs into their source or object code. To overcome this issue, developers often allow their script to be executed by any user rather than only by the user the developer originally intended to execute the script. This mismatch allows an interpreter executable to run the script with the privileges of any user.

(52) An embodiment of the method 550 ensures that the interpreter runs with only those privileges with which the parent process also runs 552. Any attempt to change the privilege level may trigger the process spawned by the interpreter to be blocked according to the IACP. Inherently this policy is a deny policy.

(53) IT Support Allow Policy and Rules

(54) Information technology (IT) personnel of an enterprise may have special scripts with specified access privileges that they may wish to execute from time to time. These scripts may be used, for example, for maintenance purposes, and may be executed when a workload is found to be in a

maintenance window, i.e., a period of time in which maintenance may be performed with minimized disruption to hosted services. Embodiments may enable IT personnel acting as an enterprise change control administrator to add specific policies. This IACP policy is an “allow” policy by nature. In an example embodiment of the method **550**, allowing or denying execution of the code of the script is thus based on permissions associated with a user requesting the execution of the code of the script.

(55) Known Attacker Tactics, Techniques, and Procedures (TTP) Deny Policy and Rules

(56) Attack frameworks, such as MITRE's ATT&CK framework referenced hereinabove, specify common techniques used by attackers to perpetuate a kill chain. This policy checks child processes against known attacker techniques, such as establishment of a reverse channel. This policy is inherently a deny policy.

(57) Known Executable Vulnerability Deny Policy and Rules

(58) From time to time, a vulnerability is disclosed in a named application, wherein if the executable is executed with specified parameters, an attacker can gain control of the application and execute arbitrary software. An embodiment leverages functionality described in U.S. Provisional Application No. 63/190,100 (introduced and incorporated by reference hereinabove). Such functionality can associate packages installed on a workload with corresponding vulnerabilities. This detail is stored in a local third-party vulnerability database (LTPCVDB). Embodiments prevent the named application from being invoked if such vulnerabilities are discovered. This policy is inherently a deny policy.

(59) Known Library Vulnerability Deny Policy and Rules

(60) From time to time, a vulnerability is disclosed in a named library wherein if the executable loads said library, an attacker can gain control of the application and execute arbitrary software. An embodiment leverages functionality described in U.S. Provisional Application No. 63/190,100 (introduced and incorporated by reference hereinabove). Such functionality can associate packages installed on a workload with corresponding vulnerabilities. This detail is stored in a local third-party vulnerability database (LTPCVDB). Embodiments prevent the named library from being loaded if such vulnerabilities are discovered. This policy is inherently a deny policy.

(61) FIG. **6** is a schematic block diagram of a system **660** for protecting a computing workload from script-based attacks, according to embodiments. The system **660** operates across an internal network **661** of an enterprise, and also connects to various remote entities via the Internet **662**. Such remote entities include an external file reputation service (EFRS) **663**, a provider package common vulnerability database (PPKCVDB) **665**, a provider 3.sup.rd party common vulnerability database (PTPCVDB) **666**, and a provider file signature repository (PFSR) **668**. The PPKCVDB **665** and the PTPCVDB **666** may connect via provider APIs **664**. The PTPCVDB **666** and the PFSR **667** may be controlled by a provider change control administrator **668**.

(62) Continuing with respect to FIG. **6**, locally to the enterprise, internal network **661** may connect to a central management server (CMS) **669**, which, in turn, may connect to a local application control policy database (LACPLDB) **670** and a local 3.sup.rd party package common vulnerability database (LTPCVDB) **671**. The LACPLDB and the LTPCVDB may be controlled by an enterprise change control administrator **672**.

(63) The aforementioned supporting elements of system **660** may function to protect a workload **640** from script-based attacks as described herein. The workload **640** may include an endpoint process monitoring client (EPMC) **673** and an application control policy (ACP) engine **674**. The ACP engine **674** may reference a workload application control policy database (WACPDB) **675**, and the EPMC **673** may reference a local file signature repository (LFSR) **676**. The EPMC may be controlled by an endpoint process monitoring driver (EPMD) **677** installed upon an operating system (OS) **678** of the workload **640**.

(64) Intercepting Process Spawn and Library Load Events

(65) The EPMC **673** of FIG. **6** can intercept process spawn and library load events. The ACP

engine **674** receives a notification of such interception from the EPMC **673** and performs further processing as specified below. Processes whose spawn events can be intercepted may be the processes known beforehand, by embodiments, as processes to be spawned, and as such may be subject to an examination of permissions thereof to ensure that such permissions match those of a parent process thereof. Alternatively, processes whose spawn events can be intercepted may be other processes not otherwise anticipated by embodiments.

(66) ACP Engine

(67) When a new process is spawned, or a library load event occurs, the ACP engine **674** tokenizes the command line in case the event is a process spawn event. Next, the ACP engine **674** examines the tokenized command line context against the various IACP policies discussed above and saved in the WACPDB **675**. Based on the policy disposition, the ACP engine **674** returns an allow/deny decision to the EPMC **673**.

(68) For library load events, the ACP engine **674** checks if a specified library has a vulnerability, and if the user has permitted the application to continue with the vulnerability. If the user has not indicated a decision (as can happen in a first instance), the ACP engine **674** may suspend the process. In return, the customer may allow the library to load within a context of a given executable. In that case, the customer's decision is conveyed to the WACPDB **675** by toggling a state flag for a combination of the given executable and the specified library. If the state flag has been turned on, then, in the future, an instance of the specified library loading into the given executable will not be blocked.

(69) In an example embodiment, the system **660** includes the EPMC **673** configured to intercept a command to commence execution of code of a script. In the example, the system **660** further includes the ACP engine **674** configured to determine, in response to the interception, whether permissions of a process to be spawned by the script match permissions of a parent process of the script, an identifier (e.g., a hash) of an executable containing the script is included in an approved list of executables, and an identifier (e.g., a hash) of the script is included in an approved list of scripts. The ACP engine **674** may be further configured to allow or deny the execution of the code of the script based on the determining, thereby protecting the workload from script-based attacks.

Advantages

(70) Embodiments have numerous advantages over existing methods for protecting computing workloads. Unlike conventional protection schemes, embodiments can automatically determine and restrict the executables, libraries, and scripts that can load in memory. This not only prevents valid executables, libraries, and scripts from loading in an unauthorized manner, but also prevents malware from loading and executing even one instruction.

(71) Conventional solutions for executable-based attacks fail to prevent attackers from launching processes with permissions that differ from parent processes thereof, indicating that it was not, in fact, the parent process that launched the process, but likely an attacker. Embodiments can automatically prevent privilege escalation attacks by ensuring a child process runs with the same privileges as the corresponding parent process, unless specifically approved by an administrator, e.g., a change control administrator. Nevertheless, embodiments allow an enterprise's information technology support personnel to add scripts into the allowed list as they see fit.

(72) Additionally, embodiments advantageously disable standards-based known attack techniques, e.g., MITRE's ATT&CK Framework. Embodiments may also disable executables or libraries from being used in a way that a known vulnerability in the executable or library can be exploited.

(73) Computer Support

(74) FIG. 7 illustrates a computer network or similar digital processing environment in which embodiments of the present disclosure may be implemented.

(75) Client computer(s)/devices **50** and server computer(s) **60** provide processing, storage, and input/output devices executing application programs and the like. The client computer(s)/devices **50** can also be linked through communications network **70** to other computing devices, including

other client devices/processes **50** and server computer(s) **60**. The communications network **70** can be part of a remote access network, a global network (e.g., the Internet), a worldwide collection of computers, local area or wide area networks, and gateways that currently use respective protocols (TCP/IP, Bluetooth®, etc.) to communicate with one another. Other electronic device/computer network architectures are suitable.

(76) Client computers/devices **50** and/or servers **60** may be configured, alone or in combination, to implement the embodiments described herein, e.g., the method **550**, amongst other examples. The server computers **60** may not be separate server computers but part of cloud network **70**.

(77) FIG. **8** is a diagram of an example internal structure of a computer (e.g., client processor/device **50** or server computers **60**) in the computer system of FIG. **7**. Each computer **50**, **60** contains a system bus **79**, where a bus is a set of hardware lines used for data transfer among the components of a computer or processing system. The system bus **79** is essentially a shared conduit that connects different elements of a computer system (e.g., processor, disk storage, memory, input/output ports, network ports, etc.) that enables the transfer of information between the elements. Attached to the system bus **79** is an input/output (I/O) device interface **82** for connecting various input and output devices (e.g., keyboard, mouse, displays, printers, speakers, etc.) to the computer **50**, **60**. A network interface **86** allows the computer to connect to various other devices attached to a network (e.g., network **70** of FIG. **7**). Memory **90** provides volatile storage for computer software instructions **92** and data **94** used to implement an embodiment of the present disclosure (e.g., the method **550**, amongst others). Disk storage **95** provides non-volatile storage for computer software instructions **92** and data **94** used to implement an embodiment of the present disclosure. A central processor unit **84** is also attached to the system bus **79** and provides for the execution of computer instructions.

(78) Embodiments or aspects thereof may be implemented in the form of hardware including but not limited to hardware circuitry, firmware, or software. If implemented in software, the software may be stored on any non-transient computer readable medium that is configured to enable a processor to load the software or subsets of instructions thereof. The processor then executes the instructions and is configured to operate or cause an apparatus to operate in a manner as described herein.

(79) Further, hardware, firmware, software, routines, or instructions may be described herein as performing certain actions and/or functions of the data processors. However, it should be appreciated that such descriptions contained herein are merely for convenience and that such actions in fact result from computing devices, processors, controllers, or other devices executing the firmware, software, routines, instructions, etc.

(80) It should be understood that the flow diagrams, block diagrams, and network diagrams may include more or fewer elements, be arranged differently, or be represented differently. But it further should be understood that certain implementations may dictate the block and network diagrams and the number of block and network diagrams illustrating the execution of the embodiments be implemented in a particular way.

(81) Accordingly, further embodiments may also be implemented in a variety of computer architectures, physical, virtual, cloud computers, and/or some combination thereof, and, thus, the data processors described herein are intended for purposes of illustration only and not as a limitation of the embodiments.

(82) The teachings of all patents, published applications and references cited herein are incorporated by reference in their entirety.

(83) While example embodiments have been particularly shown and described, it will be understood by those skilled in the art that various changes in form and details may be made therein without departing from the scope of the embodiments encompassed by the appended claims.

Claims

1. A computer-implemented method of protecting a computing workload from script-based attacks, the method comprising: receiving, at a workload, a command to commence execution of code of a script; determining whether: (i) permissions of a user issuing the command comply with a permissions security standard; (ii) an identifier of an interpreter supporting the script is included in an approved interpreter list; (iii) an identifier of a selected parameter of the interpreter is included in an approved parameter list, wherein the selected parameter modifies operations performed by the interpreter and the identifier of the selected parameter is distinct from the identifier of the interpreter; and (iv) an identifier of the script is included in an approved list of executables; and allowing or denying the execution of the code of the script based on the determining, thereby protecting the workload from script-based attacks.
2. The method of claim 1 wherein at least one of the identifier of the interpreter, the identifier of the selected parameter, and the identifier of the script, includes a checksum.
3. The method of claim 1 further comprising: receiving an identifier of an approved interpreter, or parameter thereof, from a vendor of the approved interpreter; and adding the identifier of the approved interpreter, or parameter thereof, to the approved interpreter list, or the approved parameter list.
4. The method of claim 1 further comprising: receiving an identifier of an approved script from a vendor of the approved script; and adding the identifier of the approved script to the approved list of executables.
5. The method of claim 1 wherein allowing or denying the execution of the code of the script further comprises: denying the execution of the code of the script in response to the received command not being present in a list of commands known to be used by the workload at runtime.
6. The method of claim 1 wherein allowing or denying the execution of the code of the script is further based on determining whether the user is positively authorized to execute the script based on permissions of the user and permissions of the script.
7. The method of claim 1 wherein allowing or denying the execution of the code of the script further comprises: denying the execution of the code of the script in response to a process to be spawned by the script matching a process in a list of known malicious processes.
8. The method of claim 1 wherein allowing or denying the execution of the code of the script further comprises: denying the execution of the code of the script in response to at least one of the script, a process to be spawned by the script and a library to be loaded during the execution of the script, including a known vulnerability.
9. The method of claim 1 further comprising, in response to allowing the execution of the code of the script: intercepting a given process spawned during the execution of the code of the script; tokenizing the received command; comparing tokens of the tokenized received command with entries in an allow-list database; and based on the comparing, maintaining the execution, or suspending the execution, of the code of the script, or indicating a result of the comparing to a user seeking to control execution of the script.
10. The method of claim 1 further comprising, in response to allowing the execution of the code of the script: intercepting a library loaded during the execution of the script; and in response to the library including a known vulnerability, and a disposition of the library regarding the vulnerability having not been received from a user, suspending the execution of the code of the script, or indicating the vulnerability to a user seeking to control execution of the script.
11. A system for protecting a computing workload from script-based attacks, the system comprising: an application control policy (ACP) engine and an endpoint process monitoring client (EPMC); wherein the EPMC is configured to intercept a command to commence execution of code of a script at a workload; the ACP engine is configured to determine, in response to the interception

of the command by the EPMC, whether: (i) permissions of a user issuing the command comply with a permissions security standard; (ii) an identifier of an interpreter supporting the script is included in an approved interpreter list; (iii) an identifier of a selected parameter of the interpreter is included in an approved parameter list, wherein the selected parameter modifies operations performed by the interpreter and the identifier of the selected parameter is distinct from the identifier of the interpreter; and (iv) an identifier of the script is included in an approved list of executables; and the ACP engine is configured to allow or deny the execution of the code of the script based on the determining, thereby protecting the workload from script-based attacks.

12. The system of claim 11 wherein at least one of the identifier of the interpreter, the identifier of the selected parameter, and the identifier of the script, includes a checksum.

13. The system of claim 11 wherein at least one of the EPMC and the APC engine, is further configured to: receive an identifier of an approved interpreter, or parameter thereof, from a vendor of the approved interpreter; and add the identifier of the approved interpreter, or parameter thereof, to the approved interpreter list, or the approved parameter list.

14. The system of claim 11 wherein at least one of the EPMC and the APC engine, is further configured to: receive an identifier of an approved script from a vendor of the approved script; and add the identifier of the approved script to the approved list of executables.

15. The system of claim 11 wherein allowing or denying execution of the code of the script includes: denying the execution of the code of the script in response to the intercepted command not being present in a list of commands known to be used by the workload at runtime.

16. The system of claim 11 wherein allowing or denying execution of the code of the script includes: denying the execution of the code of the script in response to a process to be spawned by the script matching a process in a list of known malicious processes.

17. The system of claim 11 wherein allowing or denying execution of the code of the script includes: denying the execution of the code of the script in response to at least one of the script, a process to be spawned by the script, and a library to be loaded during the execution of the script, including a known vulnerability.

18. The system of claim 11 wherein: the EPMC, in response to the execution of the code of the script, is configured to intercept a given process spawned during the execution of the code of the script; and the ACP engine is configured to: (i) tokenize the received command; (ii) compare tokens of the tokenized received command with entries in an allow-list database; and (iii) based on the comparing, maintain the execution, or suspend the execution, of the code of the script, or indicate a result of the comparing to a user seeking to control execution of the script.

19. The system of claim 11 wherein: the EPMC, in response to allowing the execution of the code of the script, is configured to intercept a library loaded during the execution of the script; and the ACP engine, in response to (i) the library including a known vulnerability, and (ii) a disposition of the library regarding the vulnerability having not been received from a user, is configured to suspend the execution of the code of the script, or to indicate the vulnerability to a user seeking to control execution of the script.

20. A computer program product for protecting a computing workload from script-based attacks, the computer program product comprising: one or more non-transitory computer-readable storage devices and program instructions stored on at least one of the one or more storage devices, the program instructions, when loaded and executed by a processor, cause the processor to: receive, at a workload, a command to commence execution of code of a script; determine whether: (i) permissions of a user issuing the command comply with a permissions security standard; (ii) an identifier of an interpreter supporting the script is included in an approved interpreter list; (iii) a selected parameter of the interpreter is included in an approved parameter list, wherein the selected parameter modifies operations performed by the interpreter and the identifier of the selected parameter is distinct from the identifier of the interpreter; and (iv) an identifier of the script is

included in an approved list of executables; and allow or deny the execution of the code of the script based on the determining, thereby protecting the workload from script-based attacks.
