

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12406061
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Mitra; Abhay et al.

Malware protection for virtual machines

Abstract

A computer-implemented method at a data management system comprises: receiving, at a storage appliance from a server hosting a virtual machine, a write made to the virtual machine; computing, at the storage appliance, a fingerprint of the transmitted write; comparing, at the storage appliance, the computed fingerprint to malware fingerprints in a malware catalog; repeating the computing and comparing; and disabling the virtual machine if a number of matches from the comparing breaches a predetermined threshold over a predetermined amount of time.

Inventors: Mitra; Abhay (Santa Clara, CA), Karthik; Vijay (Santa Clara, CA), Jain; Vivek Sanjay (Palo Alto, CA), Ganguli; Avishek (San Mateo, CA), Kumar; Arohi (Palo Alto, CA), Goyal; Kushaagra (Mountain View, CA), Wong; Christopher (Mountain View, CA)

Applicant: Rubrik, Inc. (Palo Alto, CA)

Family ID: 76969334

Assignee: Rubrik, Inc. (Palo Alto, CA)

Appl. No.: 18/103389

Filed: January 30, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230177157 A1	Jun. 08, 2023

Related U.S. Application Data

continuation parent-doc US 16774661 20200128 US 11604876 child-doc US 18103389

Publication Classification

Int. Cl.: **G06F21/56** (20130101); **G06F9/455** (20180101); **G06F11/14** (20060101); **G06F21/85** (20130101); **H04L9/06** (20060101)

U.S. Cl.:

CPC **G06F21/564** (20130101); **G06F9/45558** (20130101); **G06F11/1469** (20130101); **G06F21/85** (20130101); **H04L9/0643** (20130101); G06F2009/45575 (20130101); G06F2009/45587 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8307443	12/2011	Wang et al.	N/A	N/A
8997230	12/2014	McCauley et al.	N/A	N/A
9485272	12/2015	Roundy	N/A	N/A
9740577	12/2016	Chakraborty et al.	N/A	N/A
10229269	12/2018	Patton	N/A	G06F 21/554
10536471	12/2019	Derbeko et al.	N/A	N/A
10607011	12/2019	Orhan	N/A	G06F 21/566
11604876	12/2022	Mitra et al.	N/A	N/A
11616805	12/2022	Mitra et al.	N/A	N/A
2006/0130141	12/2005	Kramer	726/23	H04L 63/1408
2006/0136720	12/2005	Armstrong	713/164	G06F 21/566
2006/0236393	12/2005	Kramer et al.	N/A	N/A
2008/0168562	12/2007	Haga et al.	N/A	N/A
2008/0253380	12/2007	Cazares	370/395.53	H04L 63/0876
2009/0007100	12/2008	Field et al.	N/A	N/A
2009/0038011	12/2008	Nadathur	N/A	N/A
2009/0089879	12/2008	Wang	726/24	G06F 21/53
2009/0282483	12/2008	Bennett	726/23	H04L 63/1416
2012/0254416	12/2011	Fake et al.	N/A	N/A
2014/0143869	12/2013	Pereira	726/23	G06F 21/564
2015/0019509	12/2014	Aronovich	707/692	G06F 3/067
2016/0048683	12/2015	Sanders	726/23	H04L 63/145
2017/0195394	12/2016	Kothari	N/A	H04L 67/02
2021/0232682	12/2020	Mitra et al.	N/A	N/A

Primary Examiner: Goodchild; William J.

Attorney, Agent or Firm: Holland & Hart LLP

Background/Summary

RELATED APPLICATIONS (1) This application is a continuation application of U.S. patent application Ser. No. 16/774,661, filed Jan. 28, 2020, and entitled “MALWARE PROTECTION FOR VIRTUAL MACHINES,” which is assigned to the assignee hereof and is incorporated by reference herein in its entirety.

TECHNICAL FIELD

(1) The present disclosure generally relates to virtual machines and more particularly, but not exclusively, to malware protection for virtual machines.

BACKGROUND

(2) Malware, or malicious software, is software intentionally designed to cause damage to a computing device or act against the interest of the user of the computing device. Malware can include spyware, ransomware, viruses, etc. A virtual machine, which is an emulation of a computing device, can also be subject to malware.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) To more easily identify the discussion of any particular element or act, the most significant digit or digits in a reference number refer to the figure (“FIG.”) number in which that element or act is first introduced.

(2) FIG. 1A depicts a networked computing environment in which the disclosed technology may be practiced, according to some example embodiments.

(3) FIG. 1B depicts a server of a networked computing environment, according to some example embodiments.

(4) FIG. 1C depicts a storage appliance of a networked computing environment, according to some example embodiments.

(5) FIG. 2 shows an example virtualized infrastructure manager.

(6) FIG. 3 shows an example method for protecting a virtual machine from malware.

(7) FIG. 4 shows an example workload management system.

(8) FIG. 5 shows an example flow diagram of restoring a virtual machine.

(9) FIG. 6 is a block diagram illustrating a representative software architecture, which may be used in conjunction with various hardware architectures herein described.

(10) FIG. 7 is a block diagram illustrating components of a machine, according to some example embodiments, able to read instructions from a machine-readable medium (e.g., a machine-readable storage medium) and perform any one or more of the methodologies discussed herein.

DETAILED DESCRIPTION

(11) The description that follows includes systems, methods, techniques, instruction sequences, and computing machine program products that embody illustrative embodiments of the disclosure. In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide an understanding of various embodiments of the inventive subject matter. It will be evident, however, to those skilled in the art, that embodiments of the inventive subject matter may be practiced without these specific details. In general, well-known instruction instances, protocols, structures, and techniques are not necessarily shown in detail.

(12) Embodiments provide real-time protection against malware for virtual machines. In an embodiment, a filter driver is installed in virtualized infrastructure managers. The filter driver streams writes done in the virtual machine in real time to a cluster. The cluster can monitor these writes and look for malicious software being downloaded in all the VMs being protected.

(13) FIG. 1A depicts one embodiment of a networked computing environment **100** in which the disclosed technology may be practiced. As depicted, the networked computing environment **100** includes a datacenter **150**, a storage appliance **140**, and a computing device **154** in communication with each other via one or more networks **180**. The networked computing environment **100** may include a plurality of computing devices interconnected through one or more networks **180**. The one or more networks **180** may allow computing devices and/or storage devices to connect to and communicate with other computing devices and/or other storage devices. In some cases, the networked computing environment **100** may include other computing devices and/or other storage devices not shown. The other computing devices may include, for example, a mobile computing device, a non-mobile computing device, a server, a work station, a laptop computer, a tablet computer, a desktop computer, or an information processing system. The other storage devices may include, for example, a storage area network storage device, a networked-attached storage device, a hard disk drive, a solid-state drive, or a data storage system.

(14) The datacenter **150** may include one or more servers, such as server **160**, in communication with one or more storage devices, such as storage device **156**. The one or more servers may also be in communication with one or more storage appliances, such as storage appliance **170**. The server **160**, storage device **156**, and storage appliance **170** may be in communication with each other via a networking fabric connecting servers and data storage units within the datacenter **150** to each other. The storage appliance **170** may include a data management system for backing up virtual machines and/or files within a virtualized infrastructure. The server **160** may be used to create and manage one or more virtual machines associated with a virtualized infrastructure.

(15) The one or more virtual machines may run various applications, such as a database application or a web server (e.g., a web server hosting an auto-parts website). The storage device **156** may include one or more hardware storage devices for storing data, such as a hard disk drive (HDD), a magnetic tape drive, a solid-state drive (SSD), a storage area network (SAN) storage device, or a networked attached storage (NAS) device. In some cases, a datacenter, such as datacenter **150**, may include thousands of servers and/or data storage devices in communication with each other. The data storage devices may comprise a tiered data storage infrastructure (or a portion of a tiered data storage infrastructure). The tiered data storage infrastructure may allow for the movement of data across different tiers of a data storage infrastructure between higher-cost, higher-performance storage devices (e.g., solid-state drives and hard disk drives) and relatively lower-cost, lower-performance storage devices (e.g., magnetic tape drives).

(16) The one or more networks **180** may include a secure network such as an enterprise private network, an unsecure network such as a wireless open network, a local area network (LAN), a wide area network (WAN), and the Internet. The one or more networks **180** may include a cellular network, a mobile network, a wireless network, or a wired network. Each network of the one or more networks **180** may include hubs, bridges, routers, switches, and wired transmission media such as a direct-wired connection. The one or more networks **180** may include an extranet or other private network for securely sharing information or providing controlled access to applications or files.

(17) A server, such as server **160**, may allow a client to download information or files (e.g., executable, text, application, audio, image, or video files) from the server **160** or to perform a search query related to particular information stored on the server **160**. In some cases, a server may act as an application server or a file server. In general, a server may refer to a hardware device that acts as the host in a client-server relationship or a software process that shares a resource with or performs work for one or more clients.

(18) One embodiment of server **160** includes a network interface **165**, processor **166**, memory **167**, disk **168**, and virtualization manager **169** all in communication with each other. Network interface **165** allows server **160** to connect to one or more networks **180**. Network interface **165** may include a wireless network interface and/or a wired network interface. Processor **166** allows server **160** to

execute computer-readable instructions stored in memory **167** in order to perform processes described herein. Processor **166** may include one or more processing units, such as one or more CPUs and/or one or more GPUs. Memory **167** may comprise one or more types of memory (e.g., RAM, SRAM, DRAM, ROM, EEPROM, Flash, etc.). Disk **168** may include a hard disk drive and/or a solid-state drive. Memory **167** and disk **168** may comprise hardware storage devices.

(19) The virtualization manager **169** may manage a virtualized infrastructure and perform management operations associated with the virtualized infrastructure. The virtualization manager **169** may manage the provisioning of virtual machines running within the virtualized infrastructure and provide an interface to computing devices interacting with the virtualized infrastructure. In one example, the virtualization manager **169** may set a virtual machine into a frozen state in response to a snapshot request made via an application programming interface (API) by a storage appliance (e.g., agent installed on the storage appliance), such as storage appliance **140** or storage appliance **170**. Setting the virtual machine into a frozen state may allow a point-in-time snapshot of the virtual machine to be stored or transferred. In one example, updates made to a virtual machine that has been set into a frozen state may be written to a separate file (e.g., an update file) while the virtual machine may be set into a read-only state to prevent modifications to the virtual disk file while the virtual machine is in the frozen state.

(20) The virtualization manager **169** may then transfer data associated with the virtual machine (e.g., an image of the virtual machine or a portion of the image of the virtual disk file associated with the state of the virtual disk at a point in time is frozen) to a storage appliance in response to a request made by the storage appliance **170**. After the data associated with the point-in-time snapshot of the virtual machine has been transferred to the storage appliance **170**, the virtual machine may be released from the frozen state (i.e., unfrozen) and the updates made to the virtual machine and stored in the separate file may be merged into the virtual disk file. The virtualization manager **169** may perform various virtual machine-related tasks, such as cloning virtual machines, creating new virtual machines, monitoring the state of virtual machines, moving virtual machines between physical hosts for load balancing purposes, and facilitating backups of virtual machines.

(21) One embodiment of storage appliance **170** includes a network interface **175**, processor **176**, memory **177**, and disk **178** all in communication with each other. Network interface **175** allows storage appliance **170** to connect to one or more networks **180**. Network interface **175** may include a wireless network interface and/or a wired network interface. Processor **176** allows storage appliance **170** to execute computer-readable instructions stored in memory **177** in order to perform processes described herein. Processor **176** may include one or more processing units, such as one or more CPUs and/or one or more GPUs. Memory **177** may comprise one or more types of memory (e.g., RAM, SRAM, DRAM, ROM, EEPROM, NOR Flash, NAND Flash, etc.). Disk **178** may include a hard disk drive and/or a solid-state drive. Memory **177** and disk **178** may comprise hardware storage devices.

(22) In one embodiment, the storage appliance **170** may include four machines. Each of the four machines may include a multi-core CPU, 64 GB of RAM, a 400 GB SSD, three 4 TB HDDs, and a network interface controller. In this case, the four machines may be in communication with the one or more networks **180** via the four network interface controllers. The four machines may comprise four nodes of a server cluster. The server cluster may comprise a set of physical machines that are connected together via a network. The server cluster may be used for storing data associated with a plurality of virtual machines, such as backup data associated with different point-in-time versions of thousands of virtual machines.

(23) The networked computing environment **100** may provide a cloud computing environment for one or more computing devices. Cloud computing may refer to Internet-based computing, wherein shared resources, software, and/or information may be provided to one or more computing devices on-demand via the Internet. The networked computing environment **100** may comprise a cloud computing environment providing Software-as-a-Service (SaaS) or Infrastructure as-a-Service

(IaaS) services. SaaS may refer to a software distribution model in which applications are hosted by a service provider and made available to end users over the Internet. In one embodiment, the networked computing environment **100** may include a virtualized infrastructure that provides software, data processing, and/or data storage services to end users accessing the services via the networked computing environment **100**. In one example, networked computing environment **100** may provide cloud-based work productivity or business-related applications to a computing device, such as computing device **154**.

(24) The storage appliance **140** may comprise a cloud-based data management system for backing up virtual machines and/or files within a virtualized infrastructure, such as virtual machines running on server **160** or files stored on server **160** (e.g., locally stored files, files stored in mounted directories), according to some example embodiments.

(25) In some cases, networked computing environment **100** may provide remote access to secure applications and files stored within datacenter **150** from a remote computing device, such as computing device **154**. The datacenter **150** may use an access control application to manage remote access to protected resources, such as protected applications, databases, or files located within the datacenter **150**. To facilitate remote access to secure applications and files, a secure network connection may be established using a virtual private network (VPN). A VPN connection may allow a remote computing device, such as computing device **154**, to securely access data from a private network (e.g., from a company file server or mail server) using an unsecure public network or the Internet. The VPN connection may require client-side software (e.g., running on the remote computing device) to establish and maintain the VPN connection. The VPN client software may provide data encryption and encapsulation prior to the transmission of secure private network traffic through the Internet.

(26) In some embodiments, the storage appliance **170** may manage the extraction and storage of virtual machine snapshots associated with different point-in-time versions of one or more virtual machines running within the datacenter **150**. A snapshot of a virtual machine may correspond with a state of the virtual machine at a particular point in time. In response to a restore command from the server **160**, the storage appliance **170** may restore a point-in-time version of a virtual machine or restore point-in-time versions of one or more files located on the virtual machine and transmit the restored data to the server **160**. In response to a mount command from the server **160**, the storage appliance **170** may allow a point-in-time version of a virtual machine to be mounted and allow the server **160** to read and/or modify data associated with the point-in-time version of the virtual machine. To improve storage density, the storage appliance **170** may deduplicate and compress data associated with different versions of a virtual machine and/or deduplicate and compress data associated with different virtual machines. To improve system performance, the storage appliance **170** may first store virtual machine snapshots received from a virtualized environment in a cache, such as a flash-based cache. The cache may also store popular data or frequently accessed data (e.g., based on a history of virtual machine restorations, incremental files associated with commonly restored virtual machine versions) and current day incremental files or incremental files corresponding with snapshots captured within the past 24 hours.

(27) An incremental file may comprise a forward incremental file or a reverse incremental file. A forward incremental file may include a set of data representing changes that have occurred since an earlier point-in-time snapshot of a virtual machine. To generate a snapshot of the virtual machine corresponding with a forward incremental file, the forward incremental file may be combined with an earlier point-in-time snapshot of the virtual machine (e.g., the forward incremental file may be combined with the last full image of the virtual machine that was captured before the forward incremental file was captured and any other forward incremental files that were captured subsequent to the last full image and prior to the forward incremental file). A reverse incremental file may include a set of data representing changes from a later point-in-time snapshot of a virtual machine. To generate a snapshot of the virtual machine corresponding with a reverse incremental

file, the reverse incremental file may be combined with a later point-in-time snapshot of the virtual machine (e.g., the reverse incremental file may be combined with the most recent snapshot of the virtual machine and any other reverse incremental files that were captured prior to the most recent snapshot and subsequent to the reverse incremental file).

(28) The storage appliance **170** may provide a user interface (e.g., a web-based interface or a graphical user interface) that displays virtual machine backup information such as identifications of the virtual machines protected and the historical versions or time machine views for each of the virtual machines protected. A time machine view of a virtual machine may include snapshots of the virtual machine over a plurality of points in time. Each snapshot may comprise the state of the virtual machine at a particular point in time. Each snapshot may correspond with a different version of the virtual machine (e.g., Version 1 of a virtual machine may correspond with the state of the virtual machine at a first point in time and Version 2 of the virtual machine may correspond with the state of the virtual machine at a second point in time subsequent to the first point in time).

(29) The user interface may enable an end user of the storage appliance **170** (e.g., a system administrator or a virtualization administrator) to select a particular version of a virtual machine to be restored or mounted. When a particular version of a virtual machine has been mounted, the particular version may be accessed by a client (e.g., a virtual machine, a physical machine, or a computing device) as if the particular version was local to the client. A mounted version of a virtual machine may correspond with a mount point directory (e.g., /snapshots/VM5Nersion23). In one example, the storage appliance **170** may run an NFS server and make the particular version (or a copy of the particular version) of the virtual machine accessible for reading and/or writing. A user (e.g., database administrator) of the storage appliance **170** may then select the particular version to be mounted and run an application (e.g., a data analytics application) using the mounted version of the virtual machine. In another example, the particular version may be mounted as an iSCSI target.

(30) In some example embodiments, the storage appliance **140** is an external network connected database appliance comprising an agent **142**, an application **144**, and a storage device **146**. In some example embodiments, the application **144** is a database application for managing a database (e.g., Oracle database management system) that can store database data locally on the storage device **146**, or on remote storage locations, such as within the datacenter **150**. The agent **142** is a remote connection system for performing snapshots of database data (e.g., application **144**), and can further implement bootstrapping, upgrade, and further include backup features to transfer data from the storage appliance **140** to datacenter **150** via networks **180**.

(31) In some example embodiments, the agent **142** can be uploaded from the datacenter **150** and installed on the storage appliance **140**. After installation on storage appliance **140**, the agent **142** can be enabled or disabled by the storage appliance **140** over time. The agent **142** may acquire one or more electronic files or snapshot information associated with the one or more electronic files from the application **144**. The snapshot information may include full and/or differential snapshot data. In one example, the one or more electronic files may comprise a database file for a database and the snapshot information may comprise a differential backup of the database file.

(32) In those embodiments in which the application **144** is a database application that manages a database, the agent **142** is configured to acquire one or more electronic files corresponding with a first point-in-time version of the database from the database application. The agent **142** can further acquire a database file for the database from the application **144** or acquire a full or differential backup of the database from the computing application **144**. The determination of whether the agent **142** acquires the database file or the full or differential backup may depend on a file size of the database file. The database file may comprise a text file or a binary file. The agent **142** may transfer one or more changed data blocks corresponding with the first point-in-time version of the database to the storage appliance **140**. The one or more changed data blocks may be identified by the agent **142** by generating and comparing fingerprints or signatures for data blocks of the database file with previously generated fingerprints or signatures associated with earlier point-in-

time versions of the database file captured prior to the first point in time. In some example embodiments, the agent **142** can perform automatic upgrades or downgrades the backup agent **142** to be in-sync with software changes to a plurality of nodes (e.g., nodes operating within storage appliance **170**).

(33) In some example embodiments, the agent **142** is further configured to interface with application **144** or storage device **146** to implement changes, such as creating directories, database instances, reads/writes, and other operations to provide database management functions between the storage appliance **140** and devices within datacenter **150**. For example, the application **144** can be a relational database management application with plugin functionality, in which third-party developed plugins or extensions can be integrated in the application **144** to perform actions, such as creation of a database instance.

(34) FIG. **1B** depicts one embodiment of server **160** in FIG. **1A**. The server **160** may comprise one server out of a plurality of servers that are networked together within a datacenter **150**. In one example, the plurality of servers may be positioned within one or more server racks within the datacenter **150**. As depicted, the server **160** includes hardware-level components and software-level components. The hardware-level components include one or more processors **182**, one or more memory **184**, and one or more disks **185**. The software-level components include a hypervisor **186**, a virtualized infrastructure manager **199**, and one or more virtual machines, such as virtual machine **198**. The hypervisor **186** may comprise a native hypervisor or a hosted hypervisor. The hypervisor **186** may provide a virtual operating platform for running one or more virtual machines, such as virtual machine **198**. Virtual machine **198** includes a plurality of virtual hardware devices including a virtual processor **192**, a virtual memory **194**, and a virtual disk **195**. The virtual disk **195** may comprise a file stored within the one or more disks **185**. In one example, a virtual machine may include a plurality of virtual disks, with each virtual disk of the plurality of virtual disks associated with a different file stored on the one or more disks **185**. Virtual machine **198** may include a guest operating system **196** that runs one or more applications, such as application **197**.

(35) The virtualized infrastructure manager **199**, which may correspond with the virtualization manager **169** in FIG. **1A**, may run on a virtual machine or natively on the server **160**. The virtualized infrastructure manager **199** may provide a centralized platform for managing a virtualized infrastructure that includes a plurality of virtual machines. The virtualized infrastructure manager **199** may manage the provisioning of virtual machines running within the virtualized infrastructure and provide an interface to computing devices interacting with the virtualized infrastructure. The virtualized infrastructure manager **199** may perform various virtualized infrastructure-related tasks, such as cloning virtual machines, creating new virtual machines (e.g., new virtual machines for new nodes of the cluster), monitoring the state of virtual machines, and facilitating backups of virtual machines.

(36) In one embodiment, the server **160** may use the virtualized infrastructure manager **199** to facilitate backups for a plurality of virtual machines (e.g., eight different virtual machines) running on the server **160**. Each virtual machine running on the server **160** may run its own guest operating system and its own set of applications. Each virtual machine running on the server **160** may store its own set of files using one or more virtual disks associated with the virtual machine (e.g., each virtual machine may include two virtual disks that are used for storing data associated with the virtual machine).

(37) In one embodiment, a data management application running on a storage appliance, such as storage appliance **140** in FIG. **1A** or storage appliance **170** in FIG. **1A**, may request a snapshot of a virtual machine running on server **160**. The snapshot of the virtual machine may be stored as one or more files, with each file associated with a virtual disk of the virtual machine. A snapshot of a virtual machine may correspond with a state of the virtual machine at a particular point in time. The particular point in time may be associated with a time stamp. In one example, a first snapshot of a virtual machine may correspond with a first state of the virtual machine (including the state of

applications and files stored on the virtual machine) at a first point in time and a second snapshot of the virtual machine may correspond with a second state of the virtual machine at a second point in time subsequent to the first point in time.

(38) In response to a request for a snapshot of a virtual machine at a particular point in time, the virtualized infrastructure manager **199** may set the virtual machine into a frozen state or store a copy of the virtual machine at the particular point in time. The virtualized infrastructure manager **199** may then transfer data associated with the virtual machine (e.g., an image of the virtual machine or a portion of the image of the virtual machine) to the storage appliance. The data associated with the virtual machine may include a set of files including a virtual disk file storing contents of a virtual disk of the virtual machine at the particular point in time and a virtual machine configuration file storing configuration settings for the virtual machine at the particular point in time. The contents of the virtual disk file may include the operating system used by the virtual machine, local applications stored on the virtual disk, and user files (e.g., images and word processing documents). In some cases, the virtualized infrastructure manager **199** may transfer a full image of the virtual machine to the storage appliance or a plurality of data blocks corresponding with the full image (e.g., to enable a full image-level backup of the virtual machine to be stored on the storage appliance). In other cases, the virtualized infrastructure manager **199** may transfer a portion of an image of the virtual machine associated with data that has changed since an earlier point in time prior to the particular point in time or since a last snapshot of the virtual machine was taken. In one example, the virtualized infrastructure manager **199** may transfer only data associated with virtual blocks stored on a virtual disk of the virtual machine that have changed since the last snapshot of the virtual machine was taken. In one embodiment, the data management application may specify a first point in time and a second point in time and the virtualized infrastructure manager **199** may output one or more virtual data blocks associated with the virtual machine that have been modified between the first point in time and the second point in time.

(39) In some embodiments, the server **160** or the hypervisor **186** may communicate with a storage appliance, such as storage appliance **140** in FIG. 1A or storage appliance **170** in FIG. 1A, using a distributed file system protocol such as Network File System (NFS) Version 3. The distributed file system protocol may allow the server **160** or the hypervisor **186** to access, read, write, or modify files stored on the storage appliance as if the files were locally stored on the server **160**. The distributed file system protocol may allow the server **160** or the hypervisor **186** to mount a directory or a portion of a file system located within the storage appliance **140**. For example, the storage appliance **140** can include a standalone host of a database, where the server **160** mounts the database directories as if the files were locally stored on server **160**. Further, the server **160** may function as a backup device for storage appliance **140** by backing up data in the mounted directories in a distributed database within the datacenter **150**, such as a cluster of nodes in the storage appliance **170**.

(40) FIG. 1C depicts one embodiment of the storage appliance **170** in FIG. 1A. The storage appliance **170** may include a plurality of physical machines that may be grouped together and presented as a single computing system. Each physical machine of the plurality of physical machines may comprise a node in a cluster (e.g., a failover cluster, a Cassandra cluster). In one example, the storage appliance **170** may be positioned within a server rack within a datacenter. As depicted, the storage appliance **170** includes hardware-level components and software-level components. The hardware-level components include one or more physical machines, such as physical machine **120** and physical machine **130**. The physical machine **120** includes a network interface **121**, processor **122**, memory **123**, and disk **124** all in communication with each other. Processor **122** allows physical machine **120** to execute computer-readable instructions stored in memory **123** to perform processes described herein. Disk **124** may include a hard disk drive and/or a solid-state drive. The physical machine **130** includes a network interface **131**, processor **132**,

memory **133**, and disk **134** all in communication with each other. Processor **132** allows physical machine **130** to execute computer-readable instructions stored in memory **133** to perform processes described herein. Disk **134** may include a hard disk drive and/or a solid-state drive. In some cases, disk **134** may include a flash-based SSD or a hybrid HDD/SSD drive. In one embodiment, the storage appliance **170** may include a plurality of physical machines arranged in a cluster (e.g., four machines in a cluster). Each of the plurality of physical machines may include a plurality of multi-core CPUs, 128 GB of RAM, a 500 GB SSD, four 4 TB HDDs, and a network interface controller. (41) In some embodiments, the plurality of physical machines may be used to implement a cluster-based network fileserver. The cluster-based network file server may neither require nor use a front-end load balancer. One issue with using a front-end load balancer to host the IP address for the cluster-based network file server and to forward requests to the nodes of the cluster-based network file server is that the front-end load balancer comprises a single point of failure for the cluster-based network file server. In some cases, the file system protocol used by a server, such as server **160** in FIG. 1A, or a hypervisor, such as hypervisor **186** in FIG. 1B, to communicate with the storage appliance **170** may not provide a failover mechanism (e.g., NFS Version 3). In the case that no failover mechanism is provided on the client side, the hypervisor may not be able to connect to a new node within a cluster in the event that the node connected to the hypervisor fails.

(42) In some embodiments, each node in a cluster may be connected to each other via a network and may be associated with one or more IP addresses (e.g., two different IP addresses may be assigned to each node). In one example, each node in the cluster may be assigned a permanent IP address and a floating IP address and may be accessed using either the permanent IP address or the floating IP address. In this case, a hypervisor, such as hypervisor **186** in FIG. 1B, may be configured with a first floating IP address associated with a first node in the cluster. The hypervisor **186** may connect to the cluster using the first floating IP address. In one example, the hypervisor **186** may communicate with the cluster using the NFS Version 3 protocol. Each node in the cluster may run a Virtual Router Redundancy Protocol (VRRP) daemon. A daemon may comprise a background process. Each VRRP daemon may include a list of all floating IP addresses available within the cluster. In the event that the first node associated with the first floating IP address fails, one of the VRRP daemons may automatically assume or pick up the first floating IP address if no other VRRP daemon has already assumed the first floating IP address. Therefore, if the first node in the cluster fails or otherwise goes down, then one of the remaining VRRP daemons running on the other nodes in the cluster may assume the first floating IP address that is used by the hypervisor **186** for communicating with the cluster.

(43) In order to determine which of the other nodes in the cluster will assume the first floating IP address, a VRRP priority may be established. In one example, given a number (N) of nodes in a cluster from node(0) to node(N-1), for a floating IP address (i), the VRRP priority of nodeG) may be $G-i$ modulo N. In another example, given a number (N) of nodes in a cluster from node(0) to node(N-1), for a floating IP address (i), the VRRP priority of nodeG) may be $(i-j)$ modulo N. In these cases, nodeG) will assume floating IP address (i) only if its VRRP priority is higher than that of any other node in the cluster that is alive and announcing itself on the network. Thus, if a node fails, then there may be a clear priority ordering for determining which other node in the cluster will take over the failed node's floating IP address.

(44) In some cases, a cluster may include a plurality of nodes and each node of the plurality of nodes may be assigned a different floating IP address. In this case, a first hypervisor may be configured with a first floating IP address associated with a first node in the cluster, a second hypervisor may be configured with a second floating IP address associated with a second node in the cluster, and a third hypervisor may be configured with a third floating IP address associated with a third node in the cluster.

(45) As depicted in FIG. 1C, the software-level components of the storage appliance **170** may include data management system **102**, a virtualization interface **104**, a distributed job scheduler

108, a distributed metadata store **110**, a distributed file system **112**, and one or more virtual machine search indexes, such as virtual machine search index **106**. In one embodiment, the software-level components of the storage appliance **170** may be run using a dedicated hardware-based appliance. In another embodiment, the software-level components of the storage appliance **170** may be run from the cloud (e.g., the software-level components may be installed on a cloud service provider).

(46) In some cases, the data storage across a plurality of nodes in a cluster (e.g., the data storage available from the one or more physical machines) may be aggregated and made available over a single file system namespace (e.g., /snap-50 shots/). A directory for each virtual machine protected using the storage appliance **170** may be created (e.g., the directory for Virtual Machine A may be /snapshots/VM_A). Snapshots and other data associated with a virtual machine may reside within the directory for the virtual machine. In one example, snapshots of a virtual machine may be stored in subdirectories of the directory (e.g., a first snapshot of Virtual Machine A may reside in /snapshots/VM_A/s1/ and a second snapshot of Virtual Machine A may reside in /snapshots/VM_A/s2/).

(47) The distributed file system **112** may present itself as a single file system, in which, as new physical machines or nodes are added to the storage appliance **170**, the cluster may automatically discover the additional nodes and automatically increase the available capacity of the file system **112** for storing files and other data. Each file stored in the distributed file system **112** may be partitioned into one or more chunks or shards. Each of the one or more chunks may be stored within the distributed file system **112** as a separate file. The files stored within the distributed file system **112** may be replicated or mirrored over a plurality of physical machines, thereby creating a load-balanced and fault-tolerant distributed file system **112**. In one example, storage appliance **170** may include ten physical machines arranged as a failover cluster and a first file corresponding with a snapshot of a virtual machine (e.g., /snapshots/VM_A/s1/s1.full) may be replicated and stored on three of the ten machines.

(48) The distributed metadata store **110** may include a distributed database management system that provides high availability without a single point of failure. In one embodiment, the distributed metadata store **110** may comprise a database, such as a distributed document-oriented database. The distributed metadata store **110** may be used as a distributed key value storage system. In one example, the distributed metadata store **110** may comprise a distributed NoSQL key value store database. In some cases, the distributed metadata store **110** may include a partitioned row store, in which rows are organized into tables or other collections of related data held within a structured format within the key value store database. A table (or a set of tables) may be used to store metadata information associated with one or more files stored within the distributed file system **112**. The metadata information may include the name of a file, a size of the file, file permissions associated with the file, when the file was last modified, and file mapping information associated with an identification of the location of the file stored within a cluster of physical machines. In one embodiment, a new file corresponding with a snapshot of a virtual machine may be stored within the distributed file system **112** and metadata associated with the new file may be stored within the distributed metadata store **110**. The distributed metadata store **110** may also be used to store a backup schedule for the virtual machine and a list of snapshots for the virtual machine that are stored using the storage appliance **170**.

(49) In some cases, the distributed metadata store **110** may be used to manage one or more versions of a virtual machine. Each version of the virtual machine may correspond with a full image snapshot of the virtual machine stored within the distributed file system **112** or an incremental snapshot of the virtual machine (e.g., a forward incremental or reverse incremental) stored within the distributed file system **112**. In one embodiment, the one or more versions of the virtual machine may correspond with a plurality of files. The plurality of files may include a single full image snapshot of the virtual machine and one or more incrementals derived from the single full image

snapshot. The single full image snapshot of the virtual machine may be stored using a first storage device of a first type (e.g., a HDD) and the one or more incrementals derived from the single full image snapshot may be stored using a second storage device of a second type (e.g., an SSD). In this case, only a single full image needs to be stored and each version of the virtual machine may be generated from the single full image or the single full image combined with a subset of the one or more incrementals. Furthermore, each version of the virtual machine may be generated by performing a sequential read from the first storage device (e.g., reading a single file from a HDD) to acquire the full image and, in parallel, performing one or more reads from the second storage device (e.g., performing fast random reads from an SSD) to acquire the one or more incrementals.

(50) The distributed job scheduler **108** may be used for scheduling backup jobs that acquire and store virtual machine snapshots for one or more virtual machines over time. The distributed job scheduler **108** may follow a backup schedule to back up an entire image of a virtual machine at a particular point in time or one or more virtual disks associated with the virtual machine at the particular point in time. In one example, the backup schedule may specify that the virtual machine be backed up at a snapshot capture frequency, such as every two hours or every 24 hours. Each backup job may be associated with one or more tasks to be performed in a sequence. Each of the one or more tasks associated with a job may be run on a particular node within a cluster. In some cases, the distributed job scheduler **108** may schedule a specific job to be run on a particular node based on data stored on the particular node. For example, the distributed job scheduler **108** may schedule a virtual machine snapshot job to be run on a node in a cluster that is used to store snapshots of the virtual machine in order to reduce network congestion.

(51) The distributed job scheduler **108** may comprise a distributed fault-tolerant job scheduler, in which jobs affected by node failures are recovered and rescheduled to be run on available nodes. In one embodiment, the distributed job scheduler **108** may be fully decentralized and implemented without the existence of a master node. The distributed job scheduler **108** may run job scheduling processes on each node in a cluster or on a plurality of nodes in the cluster. In one example, the distributed job scheduler **108** may run a first set of job scheduling processes on a first node in the cluster, a second set of job scheduling processes on a second node in the cluster, and a third set of job scheduling processes on a third node in the cluster. The first set of job scheduling processes, the second set of job scheduling processes, and the third set of job scheduling processes may store information regarding jobs, schedules, and the states of jobs using a metadata store, such as distributed metadata store **110**. In the event that the first node running the first set of job scheduling processes fails (e.g., due to a network failure or a physical machine failure), the states of the jobs managed by the first set of job scheduling processes may fail to be updated within a threshold period of time (e.g., a job may fail to be completed within 30 seconds or within minutes from being started). In response to detecting jobs that have failed to be updated within the threshold period of time, the distributed job scheduler **108** may undo and restart the failed jobs on available nodes within the cluster.

(52) The job scheduling processes running on at least a plurality of nodes in a cluster (e.g., on each available node in the cluster) may manage the scheduling and execution of a plurality of jobs. The job scheduling processes may include run processes for running jobs, cleanup processes for cleaning up failed tasks, and rollback processes for rolling-back or undoing any actions or tasks performed by failed jobs. In one embodiment, the job scheduling processes may detect that a particular task for a particular job has failed and in response may perform a cleanup process to clean up or remove the effects of the particular task and then perform a rollback process that processes one or more completed tasks for the particular job in reverse order to undo the effects of the one or more completed tasks. Once the particular job with the failed task has been undone, the job scheduling processes may restart the particular job on an available node in the cluster.

(53) The distributed job scheduler **108** may manage a job in which a series of tasks associated with the job are to be performed atomically (i.e., partial execution of the series of tasks is not permitted).

If the series of tasks cannot be completely executed or there is any failure that occurs to one of the series of tasks during execution (e.g., a hard disk associated with a physical machine fails or a network connection to the physical machine fails), then the state of a data management system may be returned to a state as if none of the series of tasks were ever performed. The series of tasks may correspond with an ordering of tasks for the series of tasks and the distributed job scheduler **108** may ensure that each task of the series of tasks is executed based on the ordering of tasks. Tasks that do not have dependencies with each other may be executed in parallel.

(54) In some cases, the distributed job scheduler **108** may schedule each task of a series of tasks to be performed on a specific node in a cluster. In other cases, the distributed job scheduler **108** may schedule a first task of the series of tasks to be performed on a first node in a cluster and a second task of the series of tasks to be performed on a second node in the cluster. In these cases, the first task may have to operate on a first set of data (e.g., a first file stored in a file system) stored on the first node and the second task may have to operate on a second set of data (e.g., metadata related to the first file that is stored in a database) stored on the second node. In some embodiments, one or more tasks associated with a job may have an affinity to a specific node in a cluster.

(55) In one example, if the one or more tasks require access to a database that has been replicated on three nodes in a cluster, then the one or more tasks may be executed on one of the three nodes. In another example, if the one or more tasks require access to multiple chunks of data associated with a virtual disk that has been replicated over four nodes in a cluster, then the one or more tasks may be executed on one of the four nodes. Thus, the distributed job scheduler **108** may assign one or more tasks associated with a job to be executed on a particular node in a cluster based on the location of data required to be accessed by the one or more tasks.

(56) In one embodiment, the distributed job scheduler **108** may manage a first job associated with capturing and storing a snapshot of a virtual machine periodically (e.g., every 30 minutes). The first job may include one or more tasks, such as communicating with a virtualized infrastructure manager, such as the virtualized infrastructure manager **199** in FIG. 1B, to create a frozen copy of the virtual machine and to transfer one or more chunks (or one or more files) associated with the frozen copy to a storage appliance, such as storage appliance **170** in FIG. 1A. The one or more tasks may also include generating metadata for the one or more chunks, storing the metadata using the distributed metadata store **110**, storing the one or more chunks within the distributed file system **112**, and communicating with the virtualized infrastructure manager **199** that the frozen copy of the virtual machine may be unfrozen or released from a frozen state. The metadata for a first chunk of the one or more chunks may include information specifying a version of the virtual machine associated with the frozen copy, a time associated with the version (e.g., the snapshot of the virtual machine was taken at 5:30 p.m. on Jun. 29, 2018), and a file path to where the first chunk is stored within the distributed file system **112** (e.g., the first chunk is located at /snapshotsNM_B/s1/s1.chunk1). The one or more tasks may also include deduplication, compression (e.g., using a lossless data compression algorithm such as LZ4 or LZ77), decompression, encryption (e.g., using a symmetric key algorithm such as Triple DES or AES-256), and decryption-related tasks.

(57) The virtualization interface **104** may provide an interface for communicating with a virtualized infrastructure manager managing a virtualization infrastructure, such as virtualized infrastructure manager **199** in FIG. 1B, and requesting data associated with virtual machine snapshots from the virtualization infrastructure. The virtualization interface **104** may communicate with the virtualized infrastructure manager using an API for accessing the virtualized infrastructure manager (e.g., to communicate a request for a snapshot of a virtual machine). In this case, storage appliance **170** may request and receive data from a virtualized infrastructure without requiring agent software to be installed or running on virtual machines within the virtualized infrastructure. The virtualization interface **104** may request data associated with virtual blocks stored on a virtual disk of the virtual machine that have changed since a last snapshot of the virtual machine was taken or since a

specified prior point in time. Therefore, in some cases, if a snapshot of a virtual machine is the first snapshot taken of the virtual machine, then a full image of the virtual machine may be transferred to the storage appliance. However, if the snapshot of the virtual machine is not the first snapshot taken of the virtual machine, then only the data blocks of the virtual machine that have changed since a prior snapshot was taken may be transferred to the storage appliance.

(58) The virtual machine search index **106** may include a list of files that have been stored using a virtual machine and a version history for each of the files in the list. Each version of a file may be mapped to the earliest point-in-time snapshot of the virtual machine that includes the version of the file or to a snapshot of the virtual machine that includes the version of the file (e.g., the latest point-in-time snapshot of the virtual machine that includes the version of the file). In one example, the virtual machine search index **106** may be used to identify a version of the virtual machine that includes a particular version of a file (e.g., a particular version of a database, a spreadsheet, or a word processing document). In some cases, each of the virtual machines that are backed up or protected using storage appliance **170** may have a corresponding virtual machine search index.

(59) In one embodiment, as each snapshot of a virtual machine is ingested, each virtual disk associated with the virtual machine is parsed in order to identify a file system type associated with the virtual disk and to extract metadata (e.g., file system metadata) for each file stored on the virtual disk. The metadata may include information for locating and retrieving each file from the virtual disk. The metadata may also include a name of a file, the size of the file, the last time at which the file was modified, and a content checksum for the file. Each file that has been added, deleted, or modified since a previous snapshot was captured may be determined using the metadata (e.g., by comparing the time at which a file was last modified with a time associated with the previous snapshot). Thus, for every file that has existed within any of the snapshots of the virtual machine, a virtual machine search index may be used to identify when the file was first created (e.g., corresponding with a first version of the file) and at what times the file was modified (e.g., corresponding with subsequent versions of the file). Each version of the file may be mapped to a particular version of the virtual machine that stores that version of the file.

(60) In some cases, if a virtual machine includes a plurality of virtual disks, then a virtual machine search index may be generated for each virtual disk of the plurality of virtual disks. For example, a first virtual machine search index may catalog and map files located on a first virtual disk of the plurality of virtual disks and a second virtual machine search index may catalog and map files located on a second virtual disk of the plurality of virtual disks. In this case, a global file catalog or a global virtual machine search index for the virtual machine may include the first virtual machine search index and the second virtual machine search index. A global file catalog may be stored for each virtual machine backed up by a storage appliance within a file system, such as distributed file system **112** in FIG. **1C**.

(61) The data management system **102** may comprise an application running on the storage appliance (e.g., storage appliance **170**) that manages and stores one or more snapshots of a virtual machine. In one example, the data management system **102** may comprise a highest-level layer in an integrated software stack running on the storage appliance. The integrated software stack may include the data management system **102**, the virtualization interface **104**, the distributed job scheduler **108**, the distributed metadata store **110**, and the distributed file system **112**.

(62) In some cases, the integrated software stack may run on other computing devices, such as a server or computing device **154** in FIG. **1A**. The data management system **102** may use the virtualization interface **104**, the distributed job scheduler **108**, the distributed metadata store **110**, and the distributed file system **112** to manage and store one or more snapshots of a virtual machine. Each snapshot of the virtual machine may correspond with a point-in-time version of the virtual machine. The data management system **102** may generate and manage a list of versions for the virtual machine. Each version of the virtual machine may map to or reference one or more chunks and/or one or more files stored within the distributed file system **112**. Combined together, the one

or more chunks and/or the one or more files stored within the distributed file system **112** may comprise a full image of the version of the virtual machine.

(63) FIG. 2 shows an example virtualized infrastructure manager **199**. In an embodiment, the manager **199** comprises a filter driver **230** and remaining components reside on a different device, e.g., the storage device **156** and/or the storage appliance **140**, in which case, the filter driver **230** transmits writes from the virtual machine to the other components. The other components include a malware cataloger **210**, a malware catalog **220**, a detection engine **240**, a timer **250**, an optional bloom filter **260**, a notifier **270** and virtual machine (VM) writes **280** that the filter driver **230** transmits to the other components residing locally or at a different device. In an embodiment, the filter driver **230** and the other components can reside on separate devices. For example, the filter driver may reside on a server hosting the virtual machine and the remaining components can be located on a separate device or a same device, but not part of the actual virtual machine. In that way, writes to the VM and the transmission of writes can be done asynchronously without effecting IO bursts to the VM. Further, the VM itself cannot be compromised by malware (e.g., malware detection software disabled) as the malware detection is done outside of the VM.

(64) The malware cataloger **210** generates a catalog from a set of known malware to store in the malware catalog **220**. Four kilobyte aligned offsets (binary and/or compressed binary) are fingerprinted. In an example, fingerprinting can be done using a fingerprinting algorithm such as SHA (e.g., SHA256) and MD5 to efficiently find similarity with malware binaries or other type of algorithm. Alternatively, other sized offsets may be used.

(65) The filter driver **230** is per disk and transmits VM writes **280** to the detection engine **240**. The VM writes are also four kilobytes in length or other length to match the length of fingerprinted binaries in the catalog **220**. The detection engine **240** then computes the fingerprint for every incoming four kilobyte aligned block. If the computed fingerprint matches a fingerprint of any of the known malwares, the computed fingerprint for the write is a candidate for a match. The set of fingerprints matched for a malware along with their counts is kept track of. In an embodiment, the detection engine **240** can comprise a neural network trained using machine learning to recognize malware using a training set of known malware. The detection engine **240** could then detect malware without using the malware catalog and/or recognize malware not yet cataloged.

(66) If no matching fingerprint for a candidate is found for the last k minutes as counted by the timer **250**, the detection engine **240** discards the candidate. If more than p % of the fingerprints of a malware have matched over a predetermined time (e.g., 50% over 60 minutes), then the malware has potentially been detected. In an embodiment, the bloom filter **260** can match computed fingerprints against a bloom filter to reduce the number of fingerprints to check against known malware.

(67) In addition to or alternatively, an embodiment can use offsets where the fingerprint occurs to increase the confidence of a match. That is, finding a malware's fingerprint subsequence in the fingerprints of the VM's writes.

(68) Once malware is detected, the notifier **270** can power off affected VMs, send a notification to a user, and/or restore a VM to a point before the malware was detected. Accordingly, embodiments enable detection of malware in real-time during the download process to a VM before the malware is installed.

(69) To ensure there is no impact to the primary VM during IO bursts, the fingerprint analysis can be done asynchronously while letting the input/output go through. In case a match is detected, there is a log of all the writes that were done, so they can be replayed to revert the VM to the state before the malware was downloaded.

(70) In a multi-disk environment (e.g., RAID configurations), the detection engine **240** can combine fingerprints computed from multiple disks together to identify malware.

(71) FIG. 3 shows an example method **300** for protecting a virtual machine from malware. The method **300** may be performed by the server **160** and/or at the storage appliance **140** and/or storage

appliance **170**. For example, the filter driver **230** may transmit copies of virtual machine writes to a device not hosting the virtual machine. Alternatively, the method **300** can be performed on the same device hosting the virtual machine but outside of the virtual machine. Then, the method **300** can be performed separately from the virtual machine without disrupting virtual machine operation during input/output bursts.

(72) Initially, the malware cataloger **210** generates (**310**) the malware catalog **220** as discussed above in conjunction with the FIG. **2**. Then, the filter driver **230** transmits (**315**) a write received at the VM to outside of the VM to the detection engine **240**. Next, the detection engine **240** computes (**320**) a fingerprint for each incoming write **280** to the VM. If (**330**) there is no match, then the detection engine **240** computes (**320**) a fingerprint for the next incoming transmitted (**315**) write **280** to the VM. In an embodiment, the detection engine **240** operates asynchronously with respect to the virtual machine. For example, the virtual machine can continue to receive input/output during the operation of method **300**. If (**330**) there is a match, then the timer **250** increases (**340**) a counter. If (**350**) the counter is less than a threshold (e.g., less than 50% of incoming writes over a predetermined time), then the detection engine **240** returns to transmitting (**315**). Otherwise, if (**350**) the counter is greater than a threshold, the notifier **270** takes action (**360**), such as stopping (disabling or powering off) the VM, notifying a user via email, pop-up, text message, recorded voice call, etc., blocking source address of the malware, and/or reverting the VM to a pre-malware stage as discussed in conjunction with FIGS. **4** and **5** by reverting to a prior version of the VM using a snapshot. In addition, optionally, the bloom filter **260** can be used to expedite malware detection. Alternatively, or in addition, the detection engine **240** may comprise a neural network using machine learning to analyze the fingerprints to recognize malware without the need to access the catalog. In an embodiment, if a counter breaches a threshold then action (**360**) is taken.

(73) FIG. **4** illustrates an example of a workload management system (WMS) **400**. In one embodiment, the workload management system **400** may manage the extraction and storage of virtual machine snapshots captured at different time points of one or more virtual machines running in the virtualization manager **169** within the local datacenter **150**. In one embodiment, the workload management system **400** may manage the snapshot capturing schedule of one or more virtual machines running on one or more cloud servers. In one embodiment, the workload management system **400** may manage the extraction and storage of virtual machine snapshots captured at different time points of one or more virtual machines running in the virtualization manager **169** within the local datacenter **150** and the snapshot capturing schedule of one or more virtual machines running on one or more cloud servers. In response to a restore or recover command from a WMS client, the workload management system **400** may restore a point-in-time version of a virtual machine or restore point-in-time versions of one or more files located on the virtual machine.

(74) Referring to FIG. **4**, the workload management system **400** may have several software-level components. The software-level components of the workload management system **400** may include a cloud snapshot metadata manager **442**, an SLA policy engine **452**, a data management system **454**, a distributed job scheduler **456**, a distributed metadata store **460**, and a distributed file system **458**. In one embodiment, the software-level components of the workload management system **400** may be run using a dedicated hardware-based appliance with one or more processors and memory system. In another embodiment, the software-level components of the workload management system **400** may be run from the cloud (e.g., the software-level components may be installed on a cloud server's platform).

(75) The SLA policy engine **452** includes intelligence to determine the snapshot capturing schedule to meet terms of service level agreements between the workload management system **400** and the users, with specific aspects of the service, including how often to take virtual machine snapshots and how long to keep the snapshots, as agreed between the workload management system **400** and the users.

(76) The distributed file system **458** may present itself as a single file system in the workload management system **400** and is shared by one or more physical machines connected to the workload management system **400**. Each file stored in the distributed file system **458** may be partitioned into one or more chunks. Each of the one or more chunks may be stored within the distributed file system **458** as a separate file. The files stored within the distributed file system **458** may be replicated or mirrored over a plurality of physical machines, thereby creating a load-balanced and fault-tolerant distributed file system. In one example, workload management system **400** may include ten physical machines and a first file corresponding with a snapshot of a virtual machine (e.g.,/snapshots_A/s1/s1.full) may be replicated and stored on three of the ten machines.

(77) The distributed metadata store **460** may include a distributed database management system that provides high availability without a single point of failure. In one embodiment, the distributed metadata store **460** may comprise a database, such as a distributed document-oriented database. The distributed metadata store **460** may be used as a distributed key-value storage system. In one example, the distributed metadata store **460** may comprise a distributed NoSQL key-value store database. In some cases, the distributed metadata store **460** may include a partitioned row store, in which rows are organized into tables or other collections of related data held within a structured format within the key-value store database. A table (or a set of tables) may be used to store metadata information associated with one or more files stored within the distributed file system **458**. In one embodiment, a new file corresponding with a snapshot of a virtual machine may be stored within the distributed file system **458** and metadata associated with the new file may be stored within the distributed metadata store **460**.

(78) In some cases, the distributed metadata store **460** may be used to manage one or more versions of a virtual machine. Each version of the virtual machine may correspond with a full image snapshot of the virtual machine stored within the distributed file system **458** or an incremental snapshot of the virtual machine (e.g., a forward incremental or reverse incremental) stored within the distributed file system **458**. In one embodiment, the one or more versions of the virtual machine may correspond to a plurality of files. The plurality of files may include a single full image snapshot of the virtual machine and one or more incrementals derived from the single full image snapshot. The single full image snapshot of the virtual machine may be stored using a first storage device of a first type (e.g., an HDD) and the one or more incrementals derived from the single full image snapshot may be stored using a second storage device of a second type (e.g., an SSD). In this case, only a single full image needs to be stored, and each version of the virtual machine may be generated from the single full image or the single full image combined with a subset of the one or more incrementals. Furthermore, each version of the virtual machine may be generated by performing a sequential read from the first storage device (e.g., reading a single file from a HDD) to acquire the full image and, in parallel, performing one or more reads from the second storage device (e.g., performing fast random reads from an SSD) to acquire the one or more incrementals.

(79) The distributed job scheduler **456** may be used for scheduling backup jobs that acquire and store virtual machine snapshots for one or more virtual machines in the local datacenters and the cloud servers over time. The distributed job scheduler **456** may follow a backup schedule to back up an entire image of a virtual machine at a particular point in time or one or more data volumes associated with the virtual machine at the particular point in time. In one example, the backup schedule is the SLA agreement that prevails between the workload management system **400** and the users. Each of the one or more tasks associated with a job may be run on a particular processor of the workload management system **400**.

(80) The distributed job scheduler **456** may comprise a distributed fault tolerant job scheduler, in which jobs affected by processor failures are recovered and rescheduled to be run on available processors. The distributed job scheduler **456** may run job scheduling processes on each processor in a workload management system **400** or on a plurality of processors in the workload management system **400**. In one example, the distributed job scheduler **456** may run a first set of job scheduling

processes on a first processor in the workload management system **400**, a second set of job scheduling processes on a second processor in the workload management system **400**, and a third set of job scheduling processes on a third processor in the workload management system **400**. The first set of job scheduling processes, the second set of job scheduling processes, and the third set of job scheduling processes may store information regarding jobs, schedules, and the states of jobs using a metadata store, such as distributed metadata store **460**. In the event that the first processor running the first set of job scheduling processes fails (e.g., due to a network failure or a physical machine failure), the states of the jobs managed by the first set of job scheduling processes may fail to be updated within a threshold period of time (e.g., a job may fail to be completed within 30 seconds or within 3 minutes from being started). In response to detecting jobs that have failed to be updated within the threshold period of time, the distributed job scheduler **456** may undo and restart the failed jobs on available processors within the workload management system **400**.

(81) The cloud snapshot metadata manager **442** may have the capability to finding content in snapshots captured from virtual machines running on multiple cloud servers, compile a metadata file for the contents in the snapshots and forward the metadata file to the WMS clients. The cloud snapshot metadata manager **442** may request data associated with virtual blocks stored on a data volumes of the virtual machine that have changed since a last snapshot of the virtual machine was taken or since a specified prior point in time. Therefore, in some cases, if a snapshot of a virtual machine is the first snapshot taken of the virtual machine, then a full image of the virtual machine may be compiled to make a metadata file. However, if the snapshot of the virtual machine is not the first snapshot taken of the virtual machine, then only the data blocks of the virtual machine that have changed since a prior snapshot was taken may be compiled to make a metadata file.

(82) The data management system **454** may comprise an application running on the workload management system **400** that manages and stores one or more snapshots of a virtual machine in the local datacenter **150**. In one example, the data management system **454** may comprise a highest level layer in an integrated software stack running on the workload management system **400**. The integrated software stack may include the data management system **454**, the distributed job scheduler **456**, the distributed metadata store **460**, and the distributed file system **458**. In some cases, the integrated software stack may run on other computing devices, such as a server or computing device. The local workload management system **400** may use the distributed job scheduler **456**, the distributed metadata store **460**, and the distributed file system **458** to manage and store one or more snapshots of a virtual machine in the local datacenter. Each snapshot of the virtual machine may correspond to a point-in-time version of the virtual machine. The local workload management system **400** may generate and manage a list of versions for the virtual machine. Each version of the virtual machine may map to or reference one or more chunks and/or one or more files stored within the distributed file system **458**. Combined together, the one or more chunks and/or the one or more files stored within the distributed file system **458** may comprise a full image of the version of the virtual machine.

(83) FIG. 5 is an example workflow **1500** illustrating a representative method of restoring a virtual machine on the cloud server A **505** by the workload management system **400**. In some embodiments, the actions in the workflow may be performed in different orders and with different, fewer or additional actions than those illustrated in FIG. 5. Multiple actions can be combined in some implementations.

(84) FIG. 5 includes workflow **1500** that begins at S15.1 when a client requests, through the virtual machine search index **106**, to restore the virtual machine X on a cloud server A **505**. The request is sent to the cloud snapshot metadata manager **442** in the workload management system **400**.

(85) Workflow **1500** continues at S15.2 when the cloud snapshot metadata manager **442** instantiates a recovery virtual machine **1502** on the cloud server A **505**.

(86) At S15.3, the recovery virtual machine **1502** mounts a snapshot of the virtual machine X into its own data volumes. In one embodiment, the mounted snapshot may be selected by the user. In

another embodiment, the mounted snapshot may be the last saved snapshot. In one embodiment, the mounted snapshot may contain data for the corrupted data volumes. In another embodiment, the mounted snapshot may contain data for all the data volumes.

(87) At SI5.4, the recovery virtual machine **1502** shuts down the virtual machine X.

(88) At SI5.5, the recovery virtual machine **1502** detaches the damaged data volumes from the virtual machine X. In one embodiment, the recovery virtual machine **1502** may detach one or more corrupted data volumes, while keeping the uncorrupted volumes intact. In another embodiment, the recovery virtual machine **1502** may detach all the data volumes.

(89) At SI5.6, the recovery virtual machine **1502** detaches its own data volumes and attaches the detached data volumes to the virtual machine X. In one embodiment, the recovery virtual machine **1502** may also create a new root volume for the virtual machine X. In one embodiment, a root volume is created from an operating system image. In another embodiment, a root volume is created by mounting a snapshot containing an image of a previous version of a root volume of the virtual machine X. In one embodiment, data volumes which were mounted from snapshots are restored, while the rest of the data volumes of the virtual machine remain intact. In one embodiment, both the corrupted data volumes and the uncorrupted data volumes are restored.

(90) At SI5.7, the recovery virtual machine **1502** starts the virtual machine X with restored data volumes and root volume.

(91) FIG. **6** is a block diagram illustrating an example software architecture **606**, which may be used in conjunction with various hardware architectures herein described. FIG. **6** is a non-limiting example of a software architecture, and it will be appreciated that many other architectures may be implemented to facilitate the functionality described herein. The software architecture **606** may execute on hardware such as a machine **700** of FIG. **7** that includes, among other things, processors, memory, and I/O components. A representative hardware layer **652** is illustrated and can represent, for example, the machine **700** of FIG. **7**. The representative hardware layer **652** includes a processing unit **654** having associated executable instructions **604**. The executable instructions **604** represent the executable instructions of the software architecture **606**, including implementation of the methods, components, and so forth described herein. The hardware layer **652** also includes a memory/storage **656**, which also has the executable instructions **604**. The hardware layer **652** may also comprise other hardware **658**.

(92) In the example architecture of FIG. **6**, the software architecture **606** may be conceptualized as a stack of layers where each layer provides particular functionality. For example, the software architecture **606** may include layers such as an operating system **602**, libraries **620**, frameworks/middleware **618**, applications **616**, and a presentation layer **614**. Operationally, the applications **616** and/or other components within the layers may invoke API calls **608** through the software stack and receive a response in the form of messages **612**. The layers illustrated are representative in nature and not all software architectures have all layers. For example, some mobile or special-purpose operating systems may not provide a frameworks/middleware **618**, while others may provide such a layer. Other software architectures may include additional or different layers.

(93) The operating system **602** may manage hardware resources and provide common services. The operating system **602** may include, for example, a kernel **622**, services **624**, and drivers **626**. The kernel **622** may act as an abstraction layer between the hardware and the other software layers. For example, the kernel **622** may be responsible for memory management, processor management (e.g., scheduling), component management, networking, security settings, and so on. The services **624** may provide other common services for the other software layers. The drivers **626** are responsible for controlling or interfacing with the underlying hardware. For instance, the drivers **626** include display drivers, camera drivers, Bluetooth® drivers, flash memory drivers, serial communication drivers (e.g., Universal Serial Bus (USB) drivers), Wi-Fi® drivers, audio drivers, power management drivers, and so forth depending on the hardware configuration.

(94) The libraries **620** provide a common infrastructure that is used by the applications **616** and/or other components and/or layers. The libraries **620** provide functionality that allows other software components to perform tasks in an easier fashion than by interfacing directly with the underlying operating system **602** functionality (e.g., kernel **622**, services **624**, and/or drivers **626**). The libraries **620** may include system libraries **644** (e.g., C standard library) that may provide functions such as memory allocation functions, string manipulation functions, mathematical functions, and the like. In addition, the libraries **620** may include API libraries **646** such as media libraries (e.g., libraries to support presentation and manipulation of various media formats such as MPEG4, H.264, MP3, AAC, AMR, JPG, or PNG), graphics libraries (e.g., an OpenGL framework that may be used to render 2D and 3D graphic content on a display), database libraries (e.g., SQLite that may provide various relational database functions), web libraries (e.g., WebKit that may provide web browsing functionality), and the like. The libraries **620** may also include a wide variety of other libraries **648** to provide many other APIs to the applications **616** and other software components/modules.

(95) The frameworks/middleware **618** provide a higher-level common infrastructure that may be used by the applications **616** and/or other software components/modules. For example, the frameworks/middleware **618** may provide various graphic user interface (GUI) functions, high-level resource management, high-level location services, and so forth. The frameworks/middleware **618** may provide a broad spectrum of other APIs that may be utilized by the applications **616** and/or other software components/modules, some of which may be specific to a particular operating system **602** or platform.

(96) The applications **616** include built-in applications **638** and/or third-party applications **640**. Examples of representative built-in applications **638** may include, but are not limited to, a contacts application, a browser application, a book reader application, a location application, a media application, a messaging application, and/or a game application. The third-party applications **640** may include an application developed using the ANDROID™ or IOS™ software development kit (SDK) by an entity other than the vendor of the particular platform, and may be mobile software running on a mobile operating system such as IOS™, ANDROID™, WINDOWS® Phone, or other mobile operating systems. The third-party applications **640** may invoke the API calls **608** provided by the mobile operating system (such as the operating system **602**) to facilitate functionality described herein.

(97) The applications **616** may use built-in operating system functions (e.g., kernel **622**, services **624**, and/or drivers **626**), libraries **620**, and frameworks/middleware **618** to create user interfaces to interact with users of the system. Alternatively, or additionally, in some systems, interactions with a user may occur through a presentation layer, such as the presentation layer **614**. In these systems, the application/component “logic” can be separated from the aspects of the application/component that interact with a user.

(98) FIG. 7 is a block diagram illustrating components of a machine **700**, according to some example embodiments, able to read instructions from a machine-readable medium (e.g., a machine-readable storage medium) and perform any one or more of the methodologies discussed herein. Specifically, FIG. 7 shows a diagrammatic representation of the machine **700** in the example form of a computer system, within which instructions **716** (e.g., software, a program, an application, an applet, an app, or other executable code) for causing the machine **700** to perform any one or more of the methodologies discussed herein may be executed. As such, the instructions **716** may be used to implement modules or components described herein. The instructions **716** transform the general, non-programmed machine **700** into a particular machine **700** programmed to carry out the described and illustrated functions in the manner described. In alternative embodiments, the machine **700** operates as a standalone device or may be coupled (e.g., networked) to other machines. In a networked deployment, the machine **700** may operate in the capacity of a server machine or a client machine in a server-client network environment, or as a peer machine in a peer-to-peer (or distributed) network environment. The machine **700** may comprise, but not be limited

to, a server computer, a client computer, a personal computer (PC), a tablet computer, a laptop computer, a netbook, a set-top box (STB), a personal digital assistant (PDA), an entertainment media system, a cellular telephone, a smartphone, a mobile device, a wearable device (e.g., a smart watch), a smart home device (e.g., a smart appliance), other smart devices, a web appliance, a network router, a network switch, a network bridge, or any machine capable of executing the instructions **716**, sequentially or otherwise, that specify actions to be taken by the machine **700**. Further, while only a single machine **700** is illustrated, the term “machine” shall also be taken to include a collection of machines that individually or jointly execute the instructions **716** to perform any one or more of the methodologies discussed herein.

(99) The machine **700** may include processors **710**, memory/storage **730**, and I/O components **750**, which may be configured to communicate with each other such as via a bus **702**. The memory/storage **730** may include a main memory **732**, static memory **734**, and a storage unit **736**, both accessible to the processors **710** such as via the bus **702**. The storage unit **736** and main memory **732** store the instructions **716** embodying any one or more of the methodologies or functions described herein. The instructions **716** may also reside, completely or partially, within the static memory **734**, within the storage unit **736** (e.g., on machine readable-medium **738**), within at least one of the processors **710** (e.g., within the processor cache memory accessible to processors **712** or **714**), or any suitable combination thereof, during execution thereof by the machine **700**. Accordingly, the main memory **732**, static memory **734**, the storage unit **736**, and the memory of the processors **710** are examples of machine-readable media.

(100) The I/O components **750** may include a wide variety of components to receive input, provide output, produce output, transmit information, exchange information, capture measurements, and so on. The specific I/O components **750** that are included in a particular machine **700** will depend on the type of machine. For example, portable machines such as mobile phones will likely include a touch input device or other such input mechanisms, while a headless server machine will likely not include such a touch input device. It will be appreciated that the I/O components **750** may include many other components that are not shown in FIG. 7. The I/O components **750** are grouped according to functionality merely for simplifying the following discussion and the grouping is in no way limiting. In various example embodiments, the I/O components **750** may include output components **752** and input components **754**. The output components **752** may include visual components (e.g., a display such as a plasma display panel (PDP), a light-emitting diode (LED) display, a liquid-crystal display (LCD), a projector, or a cathode ray tube (CRT)), acoustic components (e.g., speakers), haptic components (e.g., a vibratory motor, resistance mechanisms), other signal generators, and so forth. The input components **754** may include alphanumeric input components (e.g., a keyboard, a touch screen configured to receive alphanumeric input, a photo-optical keyboard, or other alphanumeric input components), point-based input components (e.g., a mouse, a touchpad, a trackball, a joystick, a motion sensor, or other pointing instruments), tactile input components (e.g., a physical button, a touch screen that provides location and/or force of touches or touch gestures, or other tactile input components), audio input components (e.g., a microphone), and the like.

(101) In further example embodiments, the I/O components **750** may include biometric components **756**, motion components **758**, environment components **760**, or position components **762** among a wide array of other components. For example, the biometric components **756** may include components to detect expressions (e.g., hand expressions, facial expressions, vocal expressions, body gestures, or eye tracking), measure biosignals (e.g., blood pressure, heart rate, body temperature, perspiration, or brain waves), identify a person (e.g., voice identification, retinal identification, facial identification, fingerprint identification, or electroencephalogram-based identification), and the like. The motion components **758** may include acceleration sensor components (e.g., accelerometer), gravitation sensor components, rotation sensor components (e.g., gyroscope), and so forth. The environment components **760** may include, for example, illumination

sensor components (e.g., photometer), temperature sensor components (e.g., one or more thermometers that detect ambient temperature), humidity sensor components, pressure sensor components (e.g., barometer), acoustic sensor components (e.g., one or more microphones that detect background noise), proximity sensor components (e.g., infrared sensors that detect nearby objects), or other components that may provide indications, measurements, or signals corresponding to a surrounding physical environment. The position components **762** may include location sensor components (e.g., a GPS receiver component), altitude sensor components (e.g., altimeters or barometers that detect air pressure from which altitude may be derived), orientation sensor components (e.g., magnetometers), and the like.

(102) Communication may be implemented using a wide variety of technologies. The I/O components **750** may include communication components **764** operable to couple the machine **700** to a network **780** or devices **770** via a coupling **782** and a coupling **772**, respectively. For example, the communication components **764** may include a network interface component or other suitable device to interface with the network **780**. In further examples, the communication components **764** may include wired communication components, wireless communication components, cellular communication components, near field communication (NFC) components, Bluetooth® components (e.g., Bluetooth® Low Energy), Wi-Fi® components, and other communication components to provide communication via other modalities. The devices **770** may be another machine or any of a wide variety of peripheral devices (e.g., a peripheral device coupled via a USB).

(103) Moreover, the communication components **764** may detect identifiers or include components operable to detect identifiers. For example, the communication components **764** may include radio frequency identification (RFID) tag reader components, NFC smart tag detection components, optical reader components (e.g., an optical sensor to detect one-dimensional barcodes such as Universal Product Code (UPC) barcode, multi-dimensional barcodes such as Quick Response (QR) code, Aztec code, Data Matrix, Dataglyph, MaxiCode, PDF418, Ultra Code, UCC RSS-2D barcode, and other optical codes), or acoustic detection components (e.g., microphones to identify tagged audio signals). In addition, a variety of information may be derived via the communication components **764**, such as location via Internet Protocol (IP) geolocation, location via Wi-Fi® signal triangulation, location via detecting an NFC beacon signal that may indicate a particular location, and so forth.

(104) “CARRIER SIGNAL” in this context refers to any intangible medium that is capable of storing, encoding, or carrying instructions **716** for execution by the machine **700**, and includes digital or analog communications signals or other intangible media to facilitate communication of such instructions **716**. Instructions **716** may be transmitted or received over the network **780** using a transmission medium via a network interface device and using any one of a number of well-known transfer protocols.

(105) “CLIENT DEVICE” in this context refers to any machine **700** that interfaces to a network **780** to obtain resources from one or more server systems or other client devices. A client device may be, but is not limited to, a mobile phone, desktop computer, laptop, PDA, smartphone, tablet, ultrabook, netbook, multi-processor system, microprocessor-based or programmable consumer electronics system, game console, set-top box, or any other communication device that a user may use to access a network **780**.

(106) “COMMUNICATIONS NETWORK” in this context refers to one or more portions of a network **780** that may be an ad hoc network, an intranet, an extranet, a virtual private network (VPN), a local area network (LAN), a wireless LAN (WLAN), a wide area network (WAN), a wireless WAN (WWAN), a metropolitan area network (MAN), the Internet, a portion of the Internet, a portion of the Public Switched Telephone Network (PSTN), a plain old telephone service (POTS) network, a cellular telephone network, a wireless network, a Wi-Fi® network, another type of network, or a combination of two or more such networks. For example, a network

or a portion of a network **780** may include a wireless or cellular network and the coupling **782** may be a Code Division Multiple Access (CDMA) connection, a Global System for Mobile communications (GSM) connection, or another type of cellular or wireless coupling. In this example, the coupling may implement any of a variety of types of data transfer technology, such as Single Carrier Radio Transmission Technology (1×RTT), Evolution-Data Optimized (EVDO) technology, General Packet Radio Service (GPRS) technology, Enhanced Data rates for GSM Evolution (EDGE) technology, third Generation Partnership Project (3GPP) including 3G, fourth generation wireless (4G) networks, Universal Mobile Telecommunications System (UMTS), High-Speed Packet Access (HSPA), Worldwide Interoperability for Microwave Access (WiMAX), Long-Term Evolution (LTE) standard, others defined by various standard-setting organizations, other long-range protocols, or other data transfer technology.

(107) “MACHINE-READABLE MEDIUM” in this context refers to a component, a device, or other tangible media able to store instructions **716** and data temporarily or permanently and may include, but is not limited to, random-access memory (RAM), read-only memory (ROM), buffer memory, flash memory, optical media, magnetic media, cache memory, other types of storage (e.g., erasable programmable read-only memory (EPROM)), and/or any suitable combination thereof. The term “machine-readable medium” should be taken to include a single medium or multiple media (e.g., a centralized or distributed database, or associated caches and servers) able to store instructions **716**. The term “machine-readable medium” shall also be taken to include any medium, or combination of multiple media, that is capable of storing instructions **716** (e.g., code) for execution by a machine **700**, such that the instructions **716**, when executed by one or more processors **710** of the machine **700**, cause the machine **700** to perform any one or more of the methodologies described herein. Accordingly, a “machine-readable medium” refers to a single storage apparatus or device, as well as “cloud-based” storage systems or storage networks that include multiple storage apparatus or devices. The term “machine-readable medium” excludes signals per se.

(108) “COMPONENT” in this context refers to a device, a physical entity, or logic having boundaries defined by function or subroutine calls, branch points, APIs, or other technologies that provide for the partitioning or modularization of particular processing or control functions. Components may be combined via their interfaces with other components to carry out a machine process. A component may be a packaged functional hardware unit designed for use with other components and a part of a program that usually performs a particular function of related functions. Components may constitute either software components (e.g., code embodied on a machine-readable medium) or hardware components.

(109) A “hardware component” is a tangible unit capable of performing certain operations and may be configured or arranged in a certain physical manner. In various example embodiments, one or more computer systems (e.g., a standalone computer system, a client computer system, or a server computer system) or one or more hardware components of a computer system (e.g., a processor **712** or a group of processors **710**) may be configured by software (e.g., an application or application portion) as a hardware component that operates to perform certain operations as described herein. A hardware component may also be implemented mechanically, electronically, or any suitable combination thereof. For example, a hardware component may include dedicated circuitry or logic that is permanently configured to perform certain operations. A hardware component may be a special-purpose processor, such as a field-programmable gate array (FPGA) or an application-specific integrated circuit (ASIC). A hardware component may also include programmable logic or circuitry that is temporarily configured by software to perform certain operations. For example, a hardware component may include software executed by a general-purpose processor or other programmable processor. Once configured by such software, hardware components become specific machines (or specific components of a machine **700**) uniquely tailored to perform the configured functions and are no longer general-purpose processors **710**.

(110) It will be appreciated that the decision to implement a hardware component mechanically, in dedicated and permanently configured circuitry, or in temporarily configured circuitry (e.g., configured by software) may be driven by cost and time considerations. Accordingly, the phrase “hardware component” (or “hardware-implemented component”) should be understood to encompass a tangible entity, be that an entity that is physically constructed, permanently configured (e.g., hardwired), or temporarily configured (e.g., programmed) to operate in a certain manner or to perform certain operations described herein.

(111) Considering embodiments in which hardware components are temporarily configured (e.g., programmed), each of the hardware components need not be configured or instantiated at any one instance in time. For example, where a hardware component comprises a general-purpose processor **712** configured by software to become a special-purpose processor, the general-purpose processor **712** may be configured as respectively different special-purpose processors (e.g., comprising different hardware components) at different times. Software accordingly configures a particular processor **712** or processors **710**, for example, to constitute a particular hardware component at one instance of time and to constitute a different hardware component at a different instance of time.

(112) Hardware components can provide information to, and receive information from, other hardware components. Accordingly, the described hardware components may be regarded as being communicatively coupled. Where multiple hardware components exist contemporaneously, communications may be achieved through signal transmission (e.g., over appropriate circuits and buses) between or among two or more of the hardware components. In embodiments in which multiple hardware components are configured or instantiated at different times, communications between or among such hardware components may be achieved, for example, through the storage and retrieval of information in memory structures to which the multiple hardware components have access. For example, one hardware component may perform an operation and store the output of that operation in a memory device to which it is communicatively coupled. A further hardware component may then, at a later time, access the memory device to retrieve and process the stored output. Hardware components may also initiate communications with input or output devices, and can operate on a resource (e.g., a collection of information).

(113) The various operations of example methods described herein may be performed, at least partially, by one or more processors **710** that are temporarily configured (e.g., by software) or permanently configured to perform the relevant operations. Whether temporarily or permanently configured, such processors **710** may constitute processor-implemented components that operate to perform one or more operations or functions described herein. As used herein, “processor-implemented component” refers to a hardware component implemented using one or more processors **710**. Similarly, the methods described herein may be at least partially processor-implemented, with a particular processor **712** or processors **710** being an example of hardware. For example, at least some of the operations of a method may be performed by one or more processors **710** or processor-implemented components. Moreover, the one or more processors **710** may also operate to support performance of the relevant operations in a “cloud computing” environment or as a “software as a service” (SaaS). For example, at least some of the operations may be performed by a group of computers (as examples of machines **700** including processors **710**), with these operations being accessible via a network **780** (e.g., the Internet) and via one or more appropriate interfaces (e.g., an API). The performance of certain of the operations may be distributed among the processors **710**, not only residing within a single machine **700**, but deployed across a number of machines **700**. In some example embodiments, the processors **710** or processor-implemented components may be located in a single geographic location (e.g., within a home environment, an office environment, or a server farm). In other example embodiments, the processors **710** or processor-implemented components may be distributed across a number of geographic locations.

(114) “PROCESSOR” in this context refers to any circuit or virtual circuit (a physical circuit emulated by logic executing on an actual processor **712**) that manipulates data values according to

control signals (e.g., “commands,” “op codes,” “machine code,” etc.) and which produces corresponding output signals that are applied to operate a machine **700**. A processor may, for example, be a central processing unit (CPU), a reduced instruction set computing (RISC) processor, a complex instruction set computing (CISC) processor, a graphics processing unit (GPU), a digital signal processor (DSP), an ASIC, a radio-frequency integrated circuit (RFIC), or any combination thereof. A processor **710** may further be a multi-core processor **710** having two or more independent processors **712**, **714** (sometimes referred to as “cores”) that may execute instructions **716** contemporaneously.

(115) “TIMESTAMP” in this context refers to a sequence of characters or encoded information identifying when a certain event occurred, for example giving date and time of day, sometimes accurate to a small fraction of a second.

(116) The following examples describe various embodiments of methods, machine-readable media, and systems (e.g., machines, devices, or other apparatus) discussed herein.

EXAMPLES

(117) 1. A data management system, comprising: a storage appliance configured to store a snapshot of a virtual machine; one or more processors in communication with the storage appliance, the one or more processors configured to perform operations including: receiving, at the storage appliance from a server hosting the virtual machine, a write made to the virtual machine; computing, at the storage appliance, a fingerprint of the transmitted write; comparing, at the storage appliance, the computed fingerprint to malware fingerprints in a malware catalog; repeating the computing and comparing; and disabling the virtual machine if a number of matches from the comparing breaches a predetermined threshold over a predetermined amount of time.

(118) 2. The system of example 1, wherein the operations further include restoring the virtual machine using the snapshot stored in the storage appliance to a state before the predetermined threshold was breached.

(119) 3. The system of example 1, wherein the operations further include blocking writes from a source of the matches.

(120) 4. The system of example 1, wherein the operations further include generating the malware catalog including generating fingerprints of binaries and compressed binaries of known malware.

(121) 5. The system of example 4, wherein the fingerprints are computed at 4 kilobytes aligned offsets generated using SHA256.

(122) 6. The system of example 1, wherein the operations further include repeatedly generating snapshots of the virtual machine over time.

(123) 7. A computer-implemented method at a data management system, the method comprising: receiving, at a storage appliance from a server hosting a virtual machine, a write made to the virtual machine; computing, at the storage appliance, a fingerprint of the transmitted write; comparing, at the storage appliance, the computed fingerprint to malware fingerprints in a malware catalog; repeating the computing and comparing; and disabling the virtual machine if a number of matches from the comparing breaches a predetermined threshold over a predetermined amount of time.

(124) 8. The method of example 7, further comprising restoring the virtual machine using a snapshot stored in the storage appliance to a state before the predetermined threshold was breached.

(125) 9. The method of example 7, further comprising blocking writes from a source of the matches.

(126) 10. The method of example 7, further comprising generating the malware catalog including generating fingerprints of binaries and compressed binaries of known malware.

(127) 11. The method of example 10, wherein the fingerprints are computed at 4 kilobytes aligned offsets generated using SHA256.

(128) 12. The method of example 7, further comprising repeatedly generating snapshots of the virtual machine over time.

(129) 13. A non-transitory, machine-readable medium storing instructions which, when read by a

storage appliance, cause the storage appliance to perform operations comprising, at least: receiving, at the storage appliance from a server hosting a virtual machine, a write made to the virtual machine; computing, at the storage appliance, a fingerprint of the transmitted write; comparing, at the storage appliance, the computed fingerprint to malware fingerprints in a malware catalog; repeating the computing and comparing; and disabling the virtual machine if a number of matches from the comparing breaches a predetermined threshold over a predetermined amount of time.

(130) 14. The machine-readable medium of example 13, wherein the operations further include restoring the virtual machine using a snapshot stored in the storage appliance to a state before the predetermined threshold was breached.

(131) 15. The machine-readable medium of example 13, wherein the operations further include blocking writes from a source of malware.

(132) 16. The machine-readable medium of example 13, wherein the operations further include generating the malware catalog including generating fingerprints of binaries and compressed binaries of known malware.

(133) 17. The machine-readable medium of example 16, wherein the fingerprints are computed at 4 kilobytes aligned offsets generated using SHA256.

(134) 18. The machine-readable medium of example 13, wherein the operations further include repeatedly generating snapshots of the virtual machine over time.

Claims

1. A method, comprising: receiving, at a data management system, a plurality of writes made to a virtual machine; computing, at the data management system, a plurality of fingerprints for the plurality of writes made to the virtual machine; comparing, at the data management system, the plurality of fingerprints computed for the plurality of writes made to the virtual machine to a plurality of malware fingerprints in a malware catalog; and taking remedial action based at least in part on multiple matches between the plurality of fingerprints computed for the plurality of writes made to the virtual machine and the plurality of malware fingerprints breaching a threshold quantity of matches at the virtual machine within a threshold duration of time, the breaching of the threshold quantity of matches indicating a malware infection at the virtual machine, wherein taking the remedial action comprises: detaching one or more corrupted data volumes from the virtual machine; and attaching, using a snapshot of the virtual machine, one or more uncorrupted data volumes corresponding to the one or more corrupted data volumes from the snapshot to the virtual machine.

2. The method of claim 1, wherein taking the remedial action further comprises: outputting a warning for a user of the virtual machine of a presence of malware on the virtual machine.

3. The method of claim 1, further comprising: determining that malware is present on the virtual machine based at least in part on the multiple matches breaching the threshold quantity of matches within the threshold duration of time, wherein taking the remedial action is based at least in part on determining that the malware is present.

4. The method of claim 2, wherein taking the remedial action further comprises: determining a source of the malware based at least in part on a source address of the malware; and blocking additional writes from being made to the virtual machine by the source of the malware.

5. The method of claim 1, further comprising: generating the malware catalog based at least in part on generating fingerprints of binaries and compressed binaries of known malware.

6. The method of claim 1, wherein the virtual machine is hosted by a server that is separate from the data management system, and wherein the plurality of writes are received from the server.

7. An apparatus, comprising: at least one processor; memory coupled with the at least one processor; and instructions stored in the memory and executable by the at least one processor to cause the apparatus to: receive a plurality of writes made to a virtual machine; compute a plurality

of fingerprints for the plurality of writes made to the virtual machine; compare the plurality of fingerprints computed for the plurality of writes made to the virtual machine to a plurality of malware fingerprints in a malware catalog; and take remedial action based at least in part on multiple matches between the plurality of fingerprints computed for the plurality of writes made to the virtual machine and the plurality of malware fingerprints breaching a threshold quantity of matches at the virtual machine within a threshold duration of time, the breaching of the threshold quantity of matches indicating a malware infection at the virtual machine, wherein the instructions to take the remedial action are executable by the at least one processor to further: detach one or more corrupted data volumes from the virtual machine; and attach, using a snapshot of the virtual machine, one or more uncorrupted data volumes corresponding to the one or more corrupted data volumes from the snapshot to the virtual machine.

8. The apparatus of claim 7, wherein the instructions to take the remedial action are further executable by the at least one processor to: output a warning for a user of the virtual machine of a presence of malware on the virtual machine.

9. The apparatus of claim 7, wherein the instructions are further executable by the at least one processor to: determine that malware is present on the virtual machine based at least in part on the multiple matches breaching the threshold quantity of matches within the threshold duration of time, wherein taking the remedial action is based at least in part on determining that the malware is present.

10. The apparatus of claim 9, wherein the instructions to take the remedial action are further executable by the at least one processor to: determine a source of the malware based at least in part on a source address of the malware; and block additional writes from being made to the virtual machine by the source of the malware.

11. The apparatus of claim 7, wherein the virtual machine is hosted by a server that is separate from the apparatus, and wherein the plurality of writes are received from the server.

12. A non-transitory computer-readable medium storing code, the code comprising instructions executable by at least one processor to: receive, at a data management system, a plurality of writes made to a virtual machine; compute, at the data management system, a plurality of fingerprints for the plurality of writes made to the virtual machine; compare, at the data management system, the plurality of fingerprints computed for the plurality of writes made to the virtual machine to a plurality of malware fingerprints in a malware catalog; and take remedial action based at least in part on multiple matches between the plurality of fingerprints computed for the plurality of writes made to the virtual machine and the plurality of malware fingerprints breaching a threshold quantity of matches at the virtual machine within a threshold duration of time, the breaching of the threshold quantity of matches indicating a malware infection at the virtual machine, wherein, to take the remedial action, the instructions are executable by the at least one processor to further: detach one or more corrupted data volumes from the virtual machine; and attach, using a snapshot of the virtual machine, one or more uncorrupted data volumes corresponding to the one or more corrupted data volumes from the snapshot to the virtual machine.
