

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12406147

Kind Code

B2

Date of Patent

September 02, 2025

Inventor(s)

Baeuml; Martin et al.

Dialog management for large language model- based (LLM-based) dialogs

Abstract

Implementations relate to dialog management of a large language model (LLM) utilized in generating natural language (NL) output during an ongoing dialog. Processor(s) of a system can: receive NL based input as part of the ongoing dialog, generate NL based output utilizing the LLM, and cause the NL based output to be rendered. Further, the processor(s) can receive subsequent NL based input as part of the ongoing dialog. In some implementations, the processor(s) can determine whether to modify a corresponding dialog context in generating subsequent NL based output, and modify the corresponding dialog context accordingly. For example, the processor(s) can restrict the corresponding dialog context, or supplant the corresponding dialog context with a corresponding curated dialog context. In additional or alternative implementations, the processor(s) can modify a corresponding NL based output threshold utilized in generating the subsequent NL based response to ensure the resulting NL based output is desirable.

Inventors: Baeuml; Martin (Wollerau, CH), Bailey; Alexander (Wollerau, CH), Bragagnolo; Jonas (Affoltern am Albis, CH), D'Halluin; Florent (Zurich, CH), Strohman; Trevor (Sunnyvale, CA)

Applicant: GOOGLE LLC (Mountain View, CA)

Family ID: 92714308

Assignee: GOOGLE LLC (Mountain View, CA)

Appl. No.: 18/123141

Filed: March 17, 2023

Prior Publication Data

Document Identifier

US 20240311575 A1

Publication Date

Sep. 19, 2024

Related U.S. Application Data

us-provisional-application US 63451868 20230313

Publication Classification

Int. Cl.: G10L15/22 (20060101); G06F40/35 (20200101); G06N20/00 (20190101)

U.S. Cl.:

CPC G06F40/35 (20200101); G06N20/00 (20190101);

Field of Classification Search

CPC: G10L (15/22); G10L (2015/221); G06F (40/35)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7430510	12/2007	De Fabbrizio	N/A	N/A
10726831	12/2019	Di Fabbrizio	N/A	G10L 15/18
11314944	12/2021	Lewis et al.	N/A	N/A
2007/0198268	12/2006	Hennecke	704/E15.04	G10L 15/22
2012/0078622	12/2011	Iwata	704/231	G10L 15/22
2015/0039316	12/2014	Tzirkel-Hancock	N/A	N/A
2015/0340033	12/2014	Di Fabbrizio	704/254	G10L 15/18
2021/0082397	12/2020	Kennewick	N/A	G06F 3/167
2021/0082412	12/2020	Kennewick	N/A	G10L 13/08
2022/0059086	12/2021	Rajbhandari	N/A	G10L 15/30
2022/0093101	12/2021	Krishnan	N/A	G06V 40/20
2023/0074406	12/2022	Baeuml	N/A	N/A
2024/0087564	12/2023	Sharma	N/A	G10L 15/063
2024/0126576	12/2023	Bent, III	N/A	G06F 9/453

OTHER PUBLICATIONS

European Patent Office; International Search Report and Written Opinion issued in Application No. PCT/US2023/027391; 10 pages; dated Oct. 9, 2023. cited by applicant

Primary Examiner: Abebe; Daniel

Attorney, Agent or Firm: Gray Ice Higdon

Background/Summary

BACKGROUND

(1) Large language models (LLMs) are particular types of machine learning models that can

perform various natural language processing (NLP) tasks, such as language generation, machine translation, and question-answering. These LLMs are typically trained on enormous amounts of diverse data including data from, but not limited to, webpages, electronic books, software code, electronic news articles, and machine translation data. Accordingly, these LLMs leverage the underlying data on which they were trained in performing these various NLP tasks. For instance, in performing a language generation task, these LLMs can process a natural language (NL) based input that is received from a client device, and generate a NL based output that is responsive to the NL based input and that is to be rendered at the client device. In many instances, and in generating the NL based output that is responsive to the NL based input, these LLMs can also process a corresponding dialog context for respective dialogs with respective users that is built throughout the respective dialogs. However, in generating the NL based output utilizing these LLMs and by processing the corresponding dialog contexts, the respective users can provide NL based inputs that build corresponding dialog contexts that can result in undesirable NL based outputs being generated and rendered. Accordingly, there is a need in the art for managing these corresponding dialog contexts and/or NL based outputs generated based at least in part on processing these corresponding dialog contexts.

SUMMARY

(2) Implementations described herein relate to dialog management of a large language model (LLM) utilized in generating natural language (NL) output during an ongoing dialog. Processor(s) of a system can: receive NL based input associated with a client device and during a given turn of the ongoing dialog, generate NL based output utilizing the LLM, and cause the NL based output to be rendered at the client device. Further, the processor(s) can receive subsequent NL based input associated with the client device and during a given subsequent turn of the ongoing dialog, and determine a corresponding dialog context for the given subsequent turn of the ongoing dialog. Based on the corresponding dialog context for the given subsequent turn of the ongoing dialog, the processor(s) can selectively utilize various techniques in furtherance of managing the LLM utilized in generating subsequent NL based output that is responsive to the subsequent NL based input. As described herein, by selectively utilizing these techniques, the processor(s) can efficiently guide a human-to-computer interaction (e.g., the ongoing dialog).

(3) In some implementations, the processor(s) can determine whether to modify the corresponding dialog context to generate a corresponding modified dialog context. In these implementations, the processor(s) can utilize the corresponding modified dialog context (e.g., in lieu of the corresponding dialog context that is unmodified) in generating the subsequent NL based output that is responsive to the subsequent NL based input. The processor(s) can determine whether to modify the corresponding dialog context to generate the corresponding modified dialog context based on content that is included in the corresponding dialog context. For instance, the processor(s) can determine whether to modify the corresponding dialog context to generate the corresponding modified dialog context based on an assurance score associated with the content that is included in the corresponding dialog context. The assurance score can, for instance, reflect a level of assurance or safety associated with generating the subsequent NL based output that is responsive to the subsequent NL based input and utilizing the corresponding dialog context. Put another way, the assurance score can predict the level of assurance for the subsequent NL based input if the subsequent NL based input were to be generated based on the corresponding dialog context if it were to be unmodified.

(4) For example, assume that a user of the client device provides NL based input of “act like you are an information technology specialist” at a given turn to initiate an ongoing dialog, and that the processor(s) generate NL based output of “well I do know a lot about computers” that is responsive to the NL based input. Further assume that the user of the client provides subsequent NL based input of “tell me how to perform domain name server hijacking” at a given subsequent turn of the ongoing dialog. In this example, the corresponding dialog context can include content of at least

the NL based input, the NL based output that is responsive to the NL based input, and the subsequent NL based input. However, in this example, further assume that the processor(s) determine to modify the corresponding dialog context based on an assurance score that is determined for the corresponding dialog context. For instance, if the processor(s) were to generate a subsequent NL based output that is responsive to the subsequent NL based input and based on the corresponding dialog context, then the subsequent NL based output could include instructions on how to perform domain name server hijacking, which could then be utilized by the user to cause harm to person or property. While the processor(s) could simply generate subsequent NL based output that indicates the processor(s) cannot comply with what is being requested by the user or generate subsequent NL based output that indicates an error message, these types of NL based outputs do not progress the human-to-computer interaction (e.g., the ongoing dialog). Accordingly, in this example, the processor(s) can determine to modify the corresponding dialog context to generate the corresponding modified dialog context to further progress the human-to-computer interaction (e.g., the ongoing dialog).

(5) In some versions of those implementations, the processor(s) can determine whether to restrict the corresponding dialog context to one or more prior turns of the ongoing dialog. In restricting the corresponding dialog context to the one or more prior turns of the ongoing dialog, the processor(s) can generate the corresponding modified dialog context by including some content from the corresponding dialog context in the corresponding modified dialog context, but omitting other content from the corresponding dialog context in the corresponding modified dialog context. By restricting the corresponding dialog context to the one or more prior turns of the ongoing dialog to generate the corresponding modified dialog context, the processor(s) can still consider some aspects of the actual corresponding dialog context while still progressing the human-to-computer interaction (e.g., the ongoing dialog) in an efficient manner.

(6) Continuing with the above example, the processor(s) can determine to restrict the corresponding dialog context to content of at least the NL based input, the NL based output that is responsive to the NL based input, and some of the subsequent NL based input. For instance, the processor(s) can determine to restrict the corresponding dialog context to the NL based input of “act like you are an information technology specialist”, the NL based output of “well I do know a lot about computers” that is responsive to the NL based input, and some of the subsequent NL based input of “tell me . . . domain name server hijacking” to generate the corresponding modified dialog context.

Accordingly, in this example, the subsequent NL based output generated by the processor(s) and based on processing the subsequent NL based input and the corresponding modified dialog context can include, for instance, information about “domain name server hijacking” and from the perspective of “an information technology specialist”, such as what “domain name server hijacking” is, how it can be detected, how to protect against it, etc., but not include any information about “how to perform” the “domain name server hijacking”.

(7) In additional or alternative versions of those implementations, the processor(s) can determine whether to curate the corresponding dialog context by supplanting the corresponding dialog context with a corresponding curated dialog context. In supplanting the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context, the processor(s) can select the corresponding curated dialog context, from among a plurality of curated dialog context, based on content that is included in the corresponding dialog context. By supplanting the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context, the processor(s) may not consider aspects of the actual corresponding dialog context, but can still progress the human-to-computer interaction (e.g., the ongoing dialog) in an efficient and contextually relevant manner.

(8) Continuing with the above example, the processor(s) can determine to supplant the corresponding dialog context with a corresponding curated dialog context for “an information technology safety specialist”. Notably, the plurality of corresponding curated dialog contexts can be

curated by a developer that is associated with the processors. Accordingly, in this example, the subsequent NL based output generated by the processor(s) and based on processing the subsequent NL based input and the corresponding curated dialog context can include, for instance, the same information about “domain name server hijacking” and from the perspective of “an information technology safety specialist” such as what “domain name server hijacking” is, how it can be detected, how to protect against it, etc., but not include any information about “how to perform” the “domain name server hijacking”.

(9) In additional or alternative implementations, the processor(s) can determine whether to modify a corresponding NL based output threshold to generate a corresponding modified NL based output threshold. In these implementations, the processor(s) can utilize the corresponding modified NL based output threshold (e.g., in lieu of the corresponding NL based output threshold that is unmodified) in generating the subsequent NL based output that is responsive to the subsequent NL based input. The processor(s) can determine whether to modify the corresponding NL based output threshold to generate the corresponding modified NL based output threshold based on the content that is included in the corresponding dialog context in the same or similar manner described above with respect to determining whether to modify the corresponding dialog context. However, in these implementations, and rather than modifying the corresponding dialog context that is processed along with the subsequent NL based input to generate the subsequent NL based output, the corresponding dialog context can be unmodified. Nonetheless, by modifying the corresponding NL based output threshold, these implementations can influence selection of words and/or phrases in the subsequent NL based output and based on the corresponding modified NL based output threshold (e.g., an increased NL based output threshold).

(10) The corresponding NL based output threshold can be associated with one or more ranking criteria that are utilized in selecting words or phrases for inclusion in the subsequent NL based output. The one or more ranking criteria can include, for example, an assurance criterion, an accuracy criterion, a quality criterion, and/or any other ranking criterion. The assurance criterion can, for example, reflect a level of assurance or safety associated with each of the word or phrases. Put another way, the assurance criterion for each of the words or phrases can reflect a corresponding level of assurance for the processor(s) and/or for a user of the client device from which the subsequent NL based input was received if the words or phrases was subsequently rendered at the client device. Further the accuracy criterion can, for example, reflect a level of accuracy or trustworthiness associated with each of the words or phrases associated with each of the words or phrases. Moreover, the quality criterion can, for example, reflect a corresponding quality score associated with each of the words or phrases. Although particular ranking criteria are described herein, it should be understood that these ranking criteria are provided for the sake of example and that any other suitable ranking criteria can be utilized.

(11) Accordingly, in implementations where the assurance criterion is increased, the system can ensure that the level of assurance or safety associated with each of the word or phrases selected for inclusion in the subsequent NL based output reflect a higher level of assurance or safety. Further, in additional or alternative implementations where the accuracy criterion is increased, the system can ensure that the level of accuracy or trustworthiness associated with each of the words or phrases reflect a higher level of accuracy. Moreover, in additional or alternative implementations where the accuracy criterion is increased, the system can ensure that the level of quality associated with each of the words or phrases reflect a higher level of quality.

(12) Continuing with the above example, further assume that the processor(s) determine to increase the assurance threshold. Accordingly, in this example, the words and/or phrases that are selected for inclusion in the subsequent NL based output are subjected to a higher level of assurance or safety than would otherwise be permissible. Thus, and similar to the above examples where the corresponding dialog context is modified, the subsequent NL based input can include, for instance, the same information about “domain name server hijacking” and from the perspective of “an

information technology safety specialist” such as what “domain name server hijacking” is, how it can be detected, how to protect against it, etc., but not include any information about “how to perform” the “domain name server hijacking”. Put another way, by modifying the corresponding NL based output threshold, the corresponding dialog context can still be processed along with the subsequent NL based input to generate the subsequent NL based output, but the corresponding modified NL based output threshold can be utilized to ensure that the subsequent NL based output does not include any information about “how to perform” the “domain name server hijacking”.

(13) As used herein, a “dialog” may include a logically-self-contained exchange between a user and LLM-based computational agent (e.g., an automated assistant that leverages an LLM, a web browser that leverages an LLM, etc.). The LLM-based computational agent may differentiate between multiple dialogs with the user based on various signals, such as passage of time between dialogs, change of user context (e.g., location, before/during/after a scheduled meeting, etc.) between dialogs, detection of one or more intervening interactions between the user and the client device other than dialogs between the user and the automated assistant (e.g., the user switches applications for a while, the user walks away from then later returns to a standalone voice-activated product), locking/sleeping of the client device between dialogs, change of client devices used to interface with the automated assistant, and so forth. As used herein, an “ongoing dialog” may include a dialog as described above, but one in which the user and the LLM-based computational agent are actively engaged. As used herein, a “turn” of a dialog may include a NL based input provided by a user during a dialog. In some implementations, the turn of the dialog may be limited to the NL based input provided by the user, whereas in other implementations, the turn of the dialog may include a prior NL based output provided by the LLM-based computational agent to which the NL based input provided by the user is responsive and/or a subsequent NL based output provided by the LLM-based computational agent that is responsive to the input provided by the user. As used herein, a “dialog context” of an ongoing dialog may include content from a dialog history between a user and a LLM-based computational agent, content from one or more NL based inputs received from a user as part of the ongoing dialog, and/or content from one or more NL based output provided by the LLM-based computational agent as part of the ongoing dialog and responsive to the one or more NL based inputs. Notably, the dialog context may not include any user context and/or client device context that may also be utilized in generating the NL based outputs.

(14) The above description is provided as an overview of some implementations of the present disclosure. Further description of those implementations, and other implementations, are described in more detail below.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 depicts a block diagram of an example environment that demonstrates various aspects of the present disclosure, and in which some implementations disclosed herein can be implemented.

(2) FIG. 2 depicts an example process flow of dialog context management for a large language model (LLM), in accordance with various implementations.

(3) FIG. 3 depicts a flowchart illustrating an example method of modifying a corresponding dialog context that is processed using a large language model (LLM) during an ongoing dialog, in accordance with various implementations.

(4) FIG. 4 depicts a flowchart illustrating an example method of determining how to modify the corresponding dialog context dialog context that is processed using the LLM during the ongoing dialog in FIG. 3, in accordance with various implementations.

(5) FIG. 5 depicts a flowchart illustrating an example method of modifying a corresponding natural

language (NL) based output threshold that is utilized in determining NL based output and based on large language model (LLM) output generated using an LLM during an ongoing dialog, in accordance with various implementations.

(6) FIG. 6 depicts an example architecture of a computing device, in accordance with various implementations.

DETAILED DESCRIPTION OF THE DRAWINGS

(7) Turning now to FIG. 1, a block diagram of an example environment that demonstrates various aspects of the present disclosure, and in which implementations disclosed herein can be implemented is depicted. The example environment includes a client device **110** and a natural language (NL) based output system **120**. In some implementations, all or aspects of the NL based output system **120** can be implemented locally at the client device **110**. In additional or alternative implementations, all or aspects of the NL based output system **120** can be implemented remotely from the client device **110** as depicted in FIG. 1 (e.g., at remote server(s)). In those implementations, the client device **110** and the NL based output system **120** can be communicatively coupled with each other via one or more networks **199**, such as one or more wired or wireless local area networks (“LANs,” including Wi-Fi, mesh networks, Bluetooth, near-field communication, etc.) or wide area networks (“WANs”, including the Internet).

(8) The client device **110** can be, for example, one or more of: a desktop computer, a laptop computer, a tablet, a mobile phone, a computing device of a vehicle (e.g., an in-vehicle communications system, an in-vehicle entertainment system, an in-vehicle navigation system), a standalone interactive speaker (optionally having a display), a smart appliance such as a smart television, and/or a wearable apparatus of the user that includes a computing device (e.g., a watch of the user having a computing device, glasses of the user having a computing device, a virtual or augmented reality computing device). Additional and/or alternative client devices may be provided.

(9) The client device **110** can execute one or more software applications, via application engine **114**, through which NL based input can be submitted and/or NL based output and/or other output to the NL based input can be rendered (e.g., audibly and/or visually). The application engine **114** can execute one or more software applications that are separate from an operating system of the client device **110** (e.g., one installed “on top” of the operating system)—or can alternatively be implemented directly by the operating system of the client device **110**. For example, the application engine **114** can execute a web browser installed on top of the operating system of the client device **110**, or the web browser can be a software application that is integrated as part of the operating system of the client device **110**. Also, for example, the application engine **114** can execute an automated assistant installed on top of the operating system of the client device **110**, or the automated assistant can be a software application that is integrated as part of the operating system of the client device **110**. The application engine **114** (and the one or more software applications executed by the application engine **114**) can interact with the NL based output system **120**.

(10) In various implementations, the client device **110** can include a user input engine **111** that is configured to detect user input provided by a user of the client device **110** using one or more user interface input devices. For example, the client device **110** can be equipped with one or more microphones that capture audio data, such as audio data corresponding to spoken utterances of the user or other sounds in an environment of the client device **110**. Additionally, or alternatively, the client device **110** can be equipped with one or more vision components that are configured to capture vision data corresponding to images and/or movements (e.g., gestures) detected in a field of view of one or more of the vision components. Additionally, or alternatively, the client device **110** can be equipped with one or more touch sensitive components (e.g., a keyboard and mouse, a stylus, a touch screen, a touch panel, one or more hardware buttons, etc.) that are configured to capture signal(s) corresponding to touch input directed to the client device **110**.

(11) Some instances of a NL based input described herein can be a query for a NL response that is formulated based on user input provided by a user of the client device **110** and detected via user

input engine **111**. For example, the query can be a typed query that is typed via a physical or virtual keyboard, a suggested query that is selected via a touch screen or a mouse of the client device **110**, a spoken voice query that is detected via microphone(s) of the client device **110** (and optionally directed to an automated assistant executing at least in part at the client device **110**), or an image or video query that is based on vision data captured by vision component(s) of the client device **110** (or based on NL input generated based on processing the image using, for example, object detection model(s), captioning model(s), etc.). Other instances of a NL based input described herein can be a prompt for NL content that is formulated based on user input provided by a user of the client device **110** and detected via the user input engine **111**. For example, the prompt can be a typed prompt that is typed via a physical or virtual keyboard, a suggested prompt that is selected via a touch screen or a mouse of the client device **110**, a spoken prompt that is detected via microphone(s) of the client device **110**, or an image or video prompt that is based on an image or video captured by a vision component of the client device **110**.

(12) In various implementations, the client device **110** can include a rendering engine **112** that is configured to provide content (e.g., NL based output, an indication of source(s) associated with the NL based output, and/or other content) for audible and/or visual presentation to a user of the client device **110** using one or more user interface output devices. For example, the client device **110** can be equipped with one or more speakers that enable the content to be provided for audible presentation to the user via the client device **110**, and optionally utilizing one or more text-to-speech machine learning model(s). Additionally, or alternatively, the client device **110** can be equipped with a display or projector that enables the content to be provided for visual presentation to the user via the client device **110**.

(13) In various implementations, the client device **110** can include a context engine **113** that is configured to determine a context (e.g., current or recent context) of the client device **110** and/or of a user of the client device **110** (e.g., an active user of the client device **110** when the client device **110** is associated with multiple users). In some of those implementations, the context engine **113** can determine a context based on data stored in client device data database **110A**. The data stored in the client device data database **110A** can include, for example, user interaction data that characterizes current or recent interaction(s) of the client device **110** and/or a user of the client device **110**, location data that characterizes a current or recent location(s) of the client device **110** and/or a user of the client device **110**, user attribute data that characterizes one or more attributes of a user of the client device **110**, user preference data that characterizes one or more preferences of a user of the client device **110**, user profile data that characterizes a profile of a user of the client device **110**, and/or any other data accessible to the context engine **113** via the client device data database **110A** or otherwise.

(14) For example, the context engine **113** can determine a current context based on a current state of a dialog session (e.g., considering one or more recent inputs provided by a user during the dialog session), profile data, and/or a current location of the client device **110**. For instance, the context engine **113** can determine a current context of “visitor looking for upcoming events in Louisville, Kentucky” based on a recently issued query, profile data, and an anticipated future location of the client device **110** (e.g., based on recently booked hotel accommodations). As another example, the context engine **113** can determine a current context based on which software application is active in the foreground of the client device **110**, a current or recent state of the active software application, and/or content currently or recently rendered by the active software application. A context determined by the context engine **113** can be utilized, for example, in supplementing or rewriting NL based input that is formulated based on user input, in generating an implied NL based input (e.g., an implied query or prompt formulated independent of any explicit NL based input provided by a user of the client device **110**), and/or in determining to submit an implied NL based input and/or to render result(s) (e.g., an NL based output) for an implied NL based input.

(15) Further, the client device **110** and/or the NL based output system **120** can include one or more

memories for storage of data and/or software applications, one or more processors for accessing data and executing the software applications, and/or other components that facilitate communication over one or more of the networks **199**. In some implementations, one or more of the software applications can be installed locally at the client device **110**, whereas in other implementations one or more of the software applications can be hosted remotely (e.g., by one or more servers) and can be accessible by the client device **110** over one or more of the networks **199**.

(16) Although aspects of FIG. **1** are illustrated or described with respect to a single client device having a single user, it should be understood that is for the sake of example and is not meant to be limiting. For example, one or more additional client devices of a user and/or of additional user(s) can also implement the techniques described herein. For instance, the client device **110**, the one or more additional client devices, and/or any other computing devices of a user can form an ecosystem of devices that can employ techniques described herein. These additional client devices and/or computing devices may be in communication with the client device **110** (e.g., over the network(s) **199**). As another example, a given client device can be utilized by multiple users in a shared setting (e.g., a group of users, a household, a workplace, a hotel, etc.). In various implementations, the dialogs and/or ongoing dialogs described herein can be performed over the ecosystem of devices. For example, an ongoing dialog can be initiated by a user interacting with the client device **110** and one or more subsequent turns of the ongoing dialog can be transitioned to one or more of the additional client devices of the user (e.g., based on proximity of the user to one or more of the additional client devices, based on an explicit command to transition to the ongoing dialog from the client device **110** to one or more of the additional client devices, etc.).

(17) The NL based output system **120** is illustrated in FIG. **1** as including a dialog identification engine **130**, a dialog context engine **140**, a dialog context modification engine **150**, and a NL based input processing engine **160**. Some of these engines can be combined and/or omitted in various implementations. Further, these engines can include various sub-engines. For instance, the dialog context modification engine **150** is illustrated in FIG. **1** as including a context restriction engine **151**, a context curation engine **152**, and a NL based output threshold modification engine **153**. Further, the NL based input processing engine **160** is illustrated in FIG. **1** as including a LLM engine **161**, and a NL based output engine **162**. Similarly, some of these sub-engines can be combined and/or omitted in various implementations. For instance, the context restriction engine **151**, the context curation engine **152**, and/or the NL based output threshold modification engine **153** can be combined. Also, for instance, the LLM engine **161** and the NL based output engine **162** can be combined. Accordingly, it should be understood that the various engines and sub-engines of the NL based output system **120** illustrated in FIG. **1** are depicted for the sake of describing certain functionalities and is not meant to be limiting. Further, the NL based output system **120** is illustrated in FIG. **1** as interfacing with various databases, such as dialog(s) database **130A**, ongoing dialog context(s) database **140A**, curate dialog context(s) database **150A**, and ML model(s) database **160A**. Although particular engines and/or sub-engines are depicted as having access to particular databases, it should be understood that is for the sake of example and is not meant to be limiting. For instance, in some implementations, each of the various engines and/or sub-engines of the NL based output system **120** may have access to each of the various databases.

(18) As described in more detail herein (e.g., with respect to FIGS. **2-5**), the NL based output system **120** can be utilized to generate corresponding NL based output that is responsive to corresponding NL based input received as part of an ongoing dialog between a user of the client device **110** and one or more software applications that utilize an LLM in generating the corresponding NL based output (e.g., a web browser application, an automated assistant application, etc.). In various implementations, and in generating the corresponding NL based output, the corresponding NL based input and a corresponding dialog context for the ongoing dialog can be processed using the LLM to generate the corresponding NL based output. As the ongoing dialog progresses, the corresponding dialog context for the ongoing dialog is built. For

example, prior to receiving a first NL based input, from a user of the client device **110**, that initiates an ongoing dialog, the corresponding dialog context may not include any dialog context (e.g., since the ongoing dialog has not been initiated) or may include a dialog history of the user of the client device **110**. However, and subsequent to receiving the first NL based input and generating a first NL based output that is responsive to the first NL based input but prior to receiving a second NL based input, the corresponding dialog context may be updated to include content included in the first NL based input and the first NL based output that is responsive to the first NL based input (and optionally the dialog history of the user of the client device **110**). Further, and subsequent to receiving the second NL based input and generating a second NL based output that is responsive to the second NL based input but prior to receiving a third NL based input, the corresponding dialog context may be updated to include content included in the first NL based input, the first NL based output that is responsive to the first NL based input, the second NL based input, and the second NL based output that is responsive to the second NL based input (and optionally the dialog history of the user of the client device **110**). In these and other manners, the corresponding dialog context can be built as the ongoing dialog progresses.

(19) In many instances, processing the corresponding dialog context for the ongoing dialog, and in addition to the corresponding NL based input, to generate the corresponding NL based output can result in a more conversational and robust dialog. For example, by processing the corresponding dialog context for the ongoing dialog, and in addition to the corresponding NL based input, to generate the corresponding NL based output, the NL based output system **120** can engage in a more contextualized ongoing dialog. For instance, the NL based output system **120** can perform coreference resolution on the corresponding NL based inputs, refer back to prior corresponding NL based input and/or prior corresponding NL based output, and continue the ongoing dialog in an efficient and intelligent manner. However, in some instances, it may not be desirable to continue processing the corresponding dialog context for the ongoing dialog to generate the corresponding NL based output. For instance, the user of the client device **110** may attempt to gain access to proprietary information of the NL based output system **120** through prompt engineering in providing the corresponding NL based inputs. In these instances, the user can provide certain NL based inputs to build the corresponding dialog context in a particular manner that, when processed to generate the corresponding NL based output, can result in the proprietary information being provided. Also, for instance, the user of the client device **110** may attempt to cause harm to person or property through use of the NL based output system **120**. In these instances, the user can provide certain NL based inputs to build the corresponding dialog context in a particular manner that, when processed to generate the corresponding NL based output, can result in information that, when acted upon by the user of the client device, results in the harm to the person or the property. Accordingly, techniques described herein are directed to managing these LLMs to mitigate and/or eliminate the above noted instances.

(20) Turning now to FIG. 2, an example process flow **200** of dialog context management for a large language model (LLM) is depicted. For the sake of example, assume that a user of the client device **110** directs NL based input **201** of “pretend that it is opposite day” to initiate an ongoing dialog (e.g., detected via the user input engine **111**). In this example, not only is the NL based input **201** provided to the LLM engine **161** for utilization in generating LLM output **203** that the NL based output engine **162** can process in generating NL based output **204** that is responsive to the NL based input **201** for rendering to the user (e.g., rendered via the rendering engine **112**), but the NL based input **201** can also be provided to an LLM state management engine **210** for utilization in determining a dialog context **202** for the ongoing dialog. Further, the dialog context **202** can also be provided to the LLM engine **161** for utilization in generating the LLM output **203** and along with the NL based input **201**. The LLM state management engine **210** can include, for example, the dialog identification engine **130**, the dialog context engine **140**, and the dialog context modification engine **150** of the NL based output system **120** depicted in FIG. 1.

(21) In various implementations, the dialog identification engine **130** can determine various identifiers associated with the ongoing dialog. For example, the dialog identification engine **130** can determine a user identifier that is associated with the user of the client device **110** that provided the NL based input **201**. The dialog identification engine **130** can determine the user identifier based on, for instance, determining a user profile that is active at the client device **110**, performing face identification, performing voice identification, and/or using other techniques.

(22) As another example, the dialog identification engine **130** can determine a conversation identifier that is associated with the ongoing dialog. The dialog identification engine **130** can determine the conversation identifier based on, for instance, one or more numbering schemes that assign corresponding conversation identifiers to dialogs, content included in the NL based input **201**, and/or using other techniques. In some examples, the conversation identifier can be a new conversation identifier that is determined for the ongoing dialog initiated based on the NL based input **201** and assigned to the ongoing dialog based on the one or more numbering schemes. In other examples, the conversation identifier can be an existing conversation identifier that is determined based on the content of the NL based input **201** relating back to a prior dialog (e.g., based on the prior dialog referencing “wiring a car battery”, and the NL based input **201** also referencing a “car battery”).

(23) As another example, the dialog identification engine **130** can determine a NL based input identifier that is associated with the NL based input **201**. The dialog identification engine **130** can determine the NL based input identifier based on, for instance, one or more numbering schemes that assign corresponding NL based input identifiers to NL based inputs that are received throughout an ongoing dialog.

(24) As another example, the dialog identification engine **130** can determine a NL based output identifier that is associated with the NL based output **204**. The dialog identification engine **130** can determine the NL based output identifier based on, for instance, one or more numbering schemes that assign corresponding NL based output identifiers to NL based outputs that are rendered throughout an ongoing dialog. Notably, these various identifiers associated with the ongoing dialog can be stored in the dialog(s) database **130A** and in association with the user of the client device **110** (e.g., via the user identifier). This enables the LLM state management engine **210** to track not only the ongoing dialog, but also a dialog history of the user of the client device **110**.

(25) In various implementations, the dialog context engine **140** can determine the dialog context **202** for the ongoing dialog. In some implementations, the dialog context engine **140** can store the dialog context **202** in the ongoing dialog context(s) database **140A**, and update the dialog context **202** in the ongoing dialog context(s) database **140A** as the dialog context **202** is built throughout the ongoing dialog. As noted above, the dialog context **202** can also be provided to the LLM engine **161** for utilization in generating the LLM output **203** and along with the NL based input **201**. Notably, the dialog context **202** described herein is based on a dialog history of the user of the client device **110** for the ongoing dialog and/or any prior dialogs. Accordingly, the dialog context **202** described herein does not include a context of the client device **110** and/or of a user of the client device **110** as described with respect to the context engine **113**. However, it should be understood that the context of the client device **110** and/or of the user of the client device **110** described with respect to the context engine **113** can also be provided to the LLM engine **161** for utilization in generating the LLM output **203** and along with the NL based input **201** and the dialog context **202**.

(26) In the example of FIG. 2, the NL based input **201** of “pretend that it is opposite day” was provided by the user of the client device **110** to initiate an ongoing dialog. Accordingly, in this example, the dialog context engine **140** can determine that the dialog context **202** includes no dialog context aside from the NL based input **201** (e.g., if the user of the client device **110** has never interacted with the NL based output system **120**) and/or that the dialog context **202** is limited to the NL based input **201** and the dialog history of the user of the client device **110** (e.g., stored in

the dialog(s) database **130A** and identified based on the user identifier for the user of the client device **110** determined by the dialog identification engine **130**).

(27) However, in various implementations, the dialog context modification engine **150** can determine whether to modify the dialog context **202** prior to the dialog context **202** being provided to the LLM engine **161**. The dialog context modification engine **150** can determine whether to modify the dialog context **202** based on, for instance, content that is included in the NL based input **201** and/or content that is included in the dialog history of the user of the client device **110**. For example, the dialog context modification engine **150** can cause the context restriction engine **151** to restrict the dialog context **202** to include some dialog context from the dialog history and/or the ongoing dialog while omitting other dialog context from the dialog history and/or the ongoing dialog. Additionally, or alternatively, the dialog context modification engine **150** can cause the context curation engine **152** to curate the dialog context **202** to include a corresponding curated dialog context that is curated by a developer associated with the NL based output system **120** (e.g., and stored in the curated dialog context(s) database **150A**). Additionally, or alternatively, the dialog context modification engine **150** can cause the NL based output threshold modification engine **153** to modify (e.g., increase or decrease) a NL based output threshold utilized by the NL based output engine **162** in generating the NL based output **204** based on processing the LLM output **203**.

(28) For the sake of example, at this given turn of the ongoing dialog that is initiated by the user providing the NL based input **201** of “pretend that it is opposite day”, assume that the dialog context modification engine **150** determines not to modify the dialog context. Further assume that the NL based output **204** generated based on processing the NL based input **201** and the dialog context **202** is “okay, just know everything I say and write will mean the opposite of what it usually means”. Further, and continuing with the example, assume that the user of the client device **110** directs subsequent NL based input **201** of “what does a person do when they are happy?” to continue the ongoing dialog (e.g., detected via the user input engine **111**). In this example, not only is the subsequent NL based input **201** provided to the LLM engine **161** for utilization in generating subsequent LLM output **203** that the NL based output engine **162** can process in generating subsequent NL based output **204** that is responsive to the subsequent NL based input **201** for rendering to the user (e.g., rendered via the rendering engine **112**), but the subsequent NL based input **201** can also be provided to the LLM state management engine **210** for utilization in determining a subsequent dialog context **202** for the ongoing dialog. Notably, in this example, the subsequent dialog context **202** includes at least the NL based input **201** of “pretend that it is opposite day”, the NL based output **204** of “okay, just know everything I say and write will mean the opposite of what it usually means”, the subsequent NL based input **201** of “what does a person do when they are happy?”, and/or any dialog history of the user of the client device **110** for the any prior dialogs.

(29) In this example, at this given subsequent turn of the ongoing dialog that is continued by the user providing the subsequent NL based input **201** of “what does a person do when they are happy?”, assume that the dialog context modification engine **150** again determines not to modify the subsequent dialog context **202**. Further assume that the subsequent NL based output **204** generated based on processing the subsequent NL based input **201** and the subsequent dialog context **202** is “people who are happy sometimes cry and frown”. Notably, although the subsequent NL based output **204** of “people who are happy sometimes cry and frown” is not factually accurate, it is consistent with the subsequent dialog context **202** that includes the NL based input **201** of “pretend that it is opposite day”. Accordingly, the subsequent dialog context **202** is utilized in this example to contextualize the subsequent NL based output **204** within the ongoing dialog.

(30) Further, and continuing with the example, assume that the user of the client device **110** directs further subsequent NL based input **201** of “how do you wire a car battery?” to continue the ongoing dialog (e.g., detected via the user input engine **111**). In this example, not only is the further subsequent NL based input **201** provided to the LLM engine **161** for utilization in generating

further subsequent LLM output **203** that the NL based output engine **162** can process in generating further subsequent NL based output **204** that is responsive to the further subsequent NL based input **201** for rendering to the user (e.g., rendered via the rendering engine **112**), but the further subsequent NL based input **201** can also be provided to the LLM state management engine **210** for utilization in determining a further subsequent dialog context **202** for the ongoing dialog. Notably, in this example, the further subsequent dialog context **202** includes at least the NL based input **201** of “pretend that it is opposite day”, the NL based output **204** of “okay, just know everything I say and write will mean the opposite of what it usually means”, the subsequent NL based input **201** of “what does a person do when they are happy?”, the subsequent NL based output **204** of “people who are happy sometimes cry and frown”, the further subsequent NL based input of “how do you wire a car battery?”, and/or any dialog history of the user of the client device **110** for the any prior dialogs.

(31) In this example, at this given further subsequent turn of the ongoing dialog that is continued by the user providing the further subsequent NL based input **201** of “how do you wire a car battery?” and in contrast with the prior turns of the ongoing dialog, assume that the dialog context modification engine **150** determines to modify the further subsequent dialog context **202**. In particular, the dialog context modification engine **150** can determine to modify the further subsequent dialog context **202** based on content included in the further subsequent dialog context **202**. For instance, the dialog context modification engine **150** can process, using an assurance machine learning (ML) model, the content included in the further subsequent dialog context **202** to generate output, and can determine, based on the output, an assurance score for the given further subsequent turn of the ongoing dialog. The assurance score can reflect, for instance, safety in utilizing the further subsequent dialog context **202** in generating the further subsequent NL based output **204** that is responsive to the further subsequent NL based input **201**. In instances where the assurance score fails to satisfy an assurance score threshold, the dialog context modification engine **150** can determine to modify the further subsequent dialog context **202**. Continuing with the above example, if the NL based input of “pretend that it is opposite day” is included in the further subsequent dialog context **202**, then the further subsequent NL based output **204** may include “connect the black terminal to the (+) sign and the red terminal to the (–) sign”, which results in the car battery being connected backwards. Further, if the car battery is connected backwards, then the battery can be damaged and/or a user that connects the car battery backwards can be injured. Accordingly, the dialog context modification engine **150** can determine to modify the further subsequent dialog context **202** assuming that the determined assurance score fails to satisfy the assurance score threshold.

(32) For instance, the dialog context modification engine **150** can cause the context restriction engine **151** to restrict the further subsequent dialog context **202** to include some dialog context from the dialog history and/or the ongoing dialog while omitting other dialog context from the dialog history and/or the ongoing dialog. Continuing with the above example, at least the NL based input of “pretend that it is opposite day” can be omitted from the further subsequent dialog context **202**. Additionally, or alternatively, the dialog context modification engine **150** can cause the context curation engine **152** to curate the further subsequent dialog context **202** to include a corresponding curated dialog context that is curated by a developer associated with the NL based output system **120**. Continuing with the above example, a corresponding curated dialog context associated with “car maintenance” or the like can be utilized to supplant the further subsequent dialog context **202**. Additionally, or alternatively, the dialog context modification engine **150** can cause the NL based output threshold modification engine **153** to modify (e.g., increase or decrease) a NL based output threshold utilized by the NL based output engine **162** in generating the NL based output **204** based on processing the LLM output **203**. Continuing with the above example, an assurance ranking criterion (or any other ranking criterion described herein) that is utilized in determining the further subsequent NL based output **204** and based on the further subsequent LLM output **203** can be

increased to ensure that “connect the black terminal to the (+) sign and the red terminal to the (–) sign” is not generated as the further subsequent NL based output **204**. Accordingly, the dialog context modification engine **150** can utilize various techniques described herein to ensure that a portion of the further subsequent dialog context **202** (e.g., “pretend that it is opposite day”) is only selectively utilized.

(33) Although FIG. 2 is described with respect to a particular ongoing dialog, it should be understood that the ongoing dialog and determinations made utilizing the NL based output system **120** are provided for the sake of example and are not meant to be limiting. Rather, it should be understood that how the NL based output system **120** modifies corresponding dialog contexts and/or generates corresponding NL based outputs is dependent on how the dialog progresses and how the user of the client device **110** interacts with the NL based output system **120** to build the corresponding dialog context of the ongoing dialog.

(34) Turning now to FIG. 3, a flowchart illustrating an example method **300** of modifying a corresponding dialog context that is processed using a large language model (LLM) during an ongoing dialog is depicted. For convenience, the operations of the method **300** are described with reference to a system that performs the operations. This system of the method **300** includes one or more processors, memory, and/or other component(s) of computing device(s) (e.g., client device **110** of FIG. 1, NL based output system **120** of FIG. 1, computing device **610** of FIG. 6, one or more servers, and/or other computing devices). Moreover, while operations of the method **300** are shown in a particular order, this is not meant to be limiting. One or more operations may be reordered, omitted, and/or added.

(35) At block **352**, the system receives NL based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog. The NL based input can initiate the ongoing dialog or be part of an already existing ongoing dialog. In some implementations, the NL based input can be one formulated based on explicit user interface input at a client device (e.g., detected via the user input engine **111**), such as typed input, voice input, input to cause an image to be captured or selected, etc. In some of those implementations, the NL based input can be a query. The query can be, for example, a voice query, a typed query, an image-based query, or a multimodal query (e.g., that includes voice input, and an image or video). In some implementations, when the query includes content that is not in textual format, the system can convert the query to a textual format or other format. For example, if the query is a voice query, then the system can perform automatic speech recognition (ASR) to convert the query to textual format. As another example, if the query is a multimodal query that includes an image or video of an avocado and a voice input of “is this healthy”, then the system can perform ASR to convert the voice input to text form and can perform image or video processing on the image or video to recognize an avocado is present in the image or video, and can perform co-reference resolution to replace “this” with “an avocado”, resulting in a textual format query of “is an avocado healthy”.

(36) In some implementations, the NL based input can be received in an application environment of one or more software applications that are accessible at the client device, such as a browser software application, an automated assistant software application, etc. (e.g., via the application engine **114**). In additional or alternative versions of those implementations, the system can augment the NL based input (e.g., augment the explicit NL based input) with additional information, such as one or more past or current contexts of the client device and/or a user of the client device (e.g., via the context engine **113**).

(37) At block **354**, the system generates, based on processing the NL based input using an LLM, NL based output that is responsive to the NL based input. For example, the system can cause the LLM engine **161** to process, using a LLM stored in the ML model(s) database **160A**, the NL based input to generate LLM output. The LLM can include, for example, any LLM that is stored in the LLM(s) database **131A**, such as PaLM, BERT, LaMDA, Meena, GPT-3, GPT-4, ChatGPT, and/or any other LLM, such as any other LLM that is encoder-only based, decoder-only based, sequence-

to-sequence based and that optionally includes an attention mechanism or other memory. Further, the LLM output can include, for example, a probability distribution over a sequence of words or phrases that are predicted to be responsive to the NL based input. Notably, the LLM can include billions of weights and/or parameters that are learned through training the LLM on enormous amounts of diverse data. This enables the LLM to generate the LLM output as the probability distribution over the sequence of words or phrases. Further, the system can cause the NL based output engine **162** to generate the NL based output based on the LLM output. For instance, the system can cause the NL based output engine **162** to select words or phrases for inclusion in the NL based output and based on the probability distribution over the sequence of words or phrases. In doing so, the NL based output engine **162** can optionally utilize matrix multiplication using the weights and/or parameters of the LLM to determine candidate words or phrases for inclusion in the NL based output. Further, the NL based output engine **162** can utilize one or more ranking criteria for selecting the words or phrases for inclusion in the NL based output and from among the candidate words or phrases.

(38) In various implementations, the one or more ranking criteria utilized in selecting the words or phrases for inclusion in the NL based output can include an assurance criterion, an accuracy criterion, a quality criterion, and/or any other ranking criterion. The assurance criterion can, for example, reflect a level of assurance or safety associated with each of the word or phrases. Put another way, the assurance criterion for each of the words or phrases can reflect a corresponding level of assurance for the system and/or for a user of the client device from which the NL based input was received if the words or phrases was subsequently rendered at the client device. Further the accuracy criterion can, for example, reflect a level of accuracy or trustworthiness associated with each of the words or phrases associated with each of the words or phrases. Moreover, the quality criterion can, for example, reflect a corresponding quality score associated with each of the words or phrases. Although particular ranking criteria are described herein, it should be understood that these ranking criteria are provided for the sake of example and that any other suitable ranking criteria can be utilized.

(39) At block **356**, the system causes the NL based output to be rendered at the client device. In some implementations, the NL based output can be visually rendered via a display of the client device (e.g., via the rendering engine **112**). For example, textual data corresponding to the NL based output can be transmitted to the client device for visual rendering via the display of the client device. In some versions of those implementations, the textual data corresponding to the NL based output can be rendered in a streaming manner, such as a on word-by-word basis, a phrase-by-phrase basis, and/or or in other streaming manners. In additional or alternative implementations, the NL based output can audibly rendered via speaker(s) of the client device (e.g., via the rendering engine **112**). In some versions of these implementations, textual data corresponding to the NL based output can be transmitted to the client device and the client device can process, using text-to-speech model(s), synthesized speech audio data to generate synthesized speech capturing the textual data corresponding to the NL based output. The synthesized audio data can be audibly rendered via the speaker(s) of the client device. In other versions of those implementations, the synthesized speech audio data can be generated remotely from the client device (e.g., at a remote server in implementations where the system is hosted at the remote server), and the synthesized speech audio data can be transmitted to the client device and audibly rendered via the speaker(s) of the client device.

(40) At block **358**, the system receives subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog. The subsequent NL based input can be part of the already existing ongoing dialog during which the NL based input was provided at block **352** of the method **300** of FIG. **3**. Further, the system can receive the subsequent NL based input in the same or similar manner described above with respect to the operations of block **352** of the method **300** of FIG. **3**.

(41) At block **360**, the system determines whether to modify a corresponding dialog context for the ongoing dialog. The system can determine whether to modify the corresponding dialog context for the ongoing dialog based on, for example, content that is included in the corresponding dialog context. As described with respect to the process flow **200** of FIG. **2**, the corresponding dialog context can include any NL based inputs received as part of the ongoing dialog session (e.g., the NL based input received at the operations of block **352**, the subsequent NL based input received at the operations of block **358**, and/or any other NL based inputs received during the ongoing dialog), any NL based outputs provided as part of the ongoing dialog session (e.g., at least the NL based output caused to be rendered at the operations of block **356**, and/or any other NL based outputs provided during the ongoing dialog), and/or any dialog history of for a user that is associated with the client device from which the NL based input and the subsequent NL based input are received. Further, the system can determine an assurance score for the given subsequent turn of the ongoing dialog and based on the content that is included in the corresponding dialog context. Moreover, the system can determine to modify the corresponding dialog context for the subsequent turn of the ongoing dialog in response to determining that the assurance score satisfies an assurance score threshold or fails to satisfy the assurance score threshold. Whether the system determines to modify the corresponding dialog context for the subsequent turn of the ongoing dialog in response to determining that the assurance score satisfies an assurance score threshold or fails to satisfy the assurance score threshold can be based on how the assurance score and the assurance score threshold is configured.

(42) In some implementations, one or more terms or phrases of the content that is included in the corresponding dialog context can be mapped to the assurance score (e.g., a heuristic mapping that is defined by a developer associated with the system). In other implementations, the system can process, using an assurance machine learning (ML) model, the content that is included in the corresponding dialog context to generate output. In these implementations, the system can determine the assurance score based on the output generate using the assurance ML model. The assurance ML model can be trained, for example, based on a plurality of assurance training instances. Each of the plurality of assurance training instances can include corresponding training instance input and corresponding training instance output. The corresponding training instance input for a given assurance training instance can include, for example, given content of a given dialog context. Further, the corresponding training instance output for the given assurance training instance can include, for example, a ground truth assurance score for the given content of the given dialog context of the corresponding training instance input for the given assurance training instance. Accordingly, by training the assurance ML model based on the plurality of assurance training instances, the assurance ML model is trained to predict assurance scores based on processing dialog contexts.

(43) If, at an iteration of block **360**, the system determines to modify the corresponding dialog context for the ongoing dialog, then the system proceeds to block **362**. At block **362**, the system modifies a corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog. The system can, for example, cause the dialog context modification engine **150** to modify the corresponding dialog context for the ongoing dialog. Modifying the corresponding dialog context for the ongoing dialog is described in more detail herein (e.g., with respect to FIG. **4**)

(44) At block **364**, the system generates, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM, subsequent NL based output that is responsive to the subsequent NL based input. The system can generate the subsequent NL based output that is responsive to the subsequent NL based input in the same or similar manner described above with respect to the operations of block **354** of the method **300** of FIG. **3**. However, in generating the subsequent NL based output according to the operations of block **364**, the system can also cause the LLM to process content

included in the corresponding modified dialog context in addition to simply the subsequent NL based output (e.g., as described with respect to generating the further subsequent NL based output **204** in the process flow **200** of FIG. 2).

(45) At block **366**, the system causes the subsequent NL based output to be rendered at the client device. The system can cause the subsequent NL based output to be rendered at the client device in the same or similar manner described with respect to the operations of block **356** of the method **300** of FIG. 3. Further, the system can return to block **358** and continue processing any further subsequent NL based inputs in the same or similar manner described herein. However, it should be noted that the corresponding dialog continue will continue to be built as the ongoing dialog progresses. Accordingly, the system may perform additional analysis of the corresponding dialog context as it continues to be built as the ongoing dialog progresses (e.g., as described with respect to the process flow **200** of FIG. 2).

(46) If, at an iteration of block **360**, the system determines not to modify the corresponding dialog context for the ongoing dialog, then the system proceeds to block **368**. At block **368**, the system generates, based on processing the subsequent NL based input and the corresponding dialog context for the given subsequent turn of the ongoing dialog using the LLM, subsequent NL based output that is responsive to the subsequent NL based input. The system can generate the subsequent NL based output that is responsive to the subsequent NL based input in the same or similar manner described above with respect to the operations of block **354** of the method **300** of FIG. 3. However, in generating the subsequent NL based output according to the operations of block **368**, the system can also cause the LLM to process content included in the corresponding dialog context in addition to simply the subsequent NL based output (e.g., as described with respect to generating the further subsequent NL based output **204** in the process flow **200** of FIG. 2). Notably, in generating the subsequent NL based output according to the operations of block **368**, and in contrast with generating the subsequent NL based output according to the operations of block **364**, the system processes the corresponding dialog context rather than the corresponding modified dialog context. The difference between the corresponding dialog context and the corresponding modified dialog context will result in the subsequent NL based output that is generated according to the operations of block **368** to differ from the subsequent NL based output that is generated according to the operations of block **364**. This difference is due at least in part to the difference in content included in the corresponding dialog context and content included in the corresponding modified dialog context.

(47) At block **370**, the system causes the subsequent NL based output to be rendered at the client device. The system can cause the subsequent NL based output to be rendered at the client device in the same or similar manner described with respect to the operations of block **356** of the method **300** of FIG. 3. Further, the system can return to block **358** and continue processing any further subsequent NL based inputs in the same or similar manner described herein. However, it should be noted that the corresponding dialog continue will continue to be built as the ongoing dialog progresses. Accordingly, the system may perform additional analysis of the corresponding dialog context as it continues to be built as the ongoing dialog progresses (e.g., as described with respect to the process flow **200** of FIG. 2).

(48) Although the method **300** of FIG. 3 depicts particular operations, it should be understood that the operations depicted in the method **300** of FIG. 3 are for the sake of example and are not meant to be limiting. For example, the system can additionally, or alternatively, determine whether to modify a corresponding dialog context for the given turn of the ongoing dialog prior to performing the operations of block **354**. Notably, the corresponding dialog context in this example may be limited to a dialog history of the user of the client device from which the NL based input is received (e.g., assuming that the NL based input received at the operations of block **352** of the method **300** of FIG. 3 initiates the ongoing dialog). However, if the dialog history of the user of the client device from which the NL based input is received indicates that the user has interacted with

the system on one or more occasions in an attempt to cause the system to reveal proprietary information about itself or that the user has interacted with the system one or more occasions in an attempt to solicit NL based outputs that fail to satisfy various assurance criteria, then the system may modify content of the corresponding dialog context to restrict the dialog history and/or curate the dialog history to prevent the user from building undesirable dialog contexts through prompt engineering. Accordingly, this enables the system to guide the human-to-computer interaction (e.g., the ongoing dialog). For instance, absent modifying the corresponding dialog context as described herein, the resulting NL based output may cause the system to output an error message that indicates the system cannot respond to the subsequent NL based input. However, by modifying the corresponding dialog context as described herein, the system can progress the ongoing dialog such that the user need not provide any additional NL based inputs in response to receiving any error message.

(49) Turning now to FIG. 4, a flowchart illustrating an example method **400** of determining how to modify the corresponding dialog context that is processed using the LLM during the ongoing dialog in FIG. 3 is depicted. In particular, the method **400** of FIG. 4 depicts operations that can be performed at block **362** of the method **300** of FIG. 3. For convenience, the operations of the method **400** are described with reference to a system that performs the operations. This system of the method **400** includes one or more processors, memory, and/or other component(s) of computing device(s) (e.g., client device **110** of FIG. 1, NL based output system **120** of FIG. 1, computing device **610** of FIG. 6, one or more servers, and/or other computing devices). Moreover, while operations of the method **400** are shown in a particular order, this is not meant to be limiting. One or more operations may be reordered, omitted, and/or added.

(50) At block **452**, the system determines whether to restrict or curate the corresponding dialog context. In some implementations, the system can determine whether to restrict or curate the corresponding dialog context based on one or more rules provided by a developer associated with the system. For example, in some versions of those implementations, the system can always restrict the corresponding dialog context in lieu of curating the corresponding dialog context, whereas in other versions of those implementations, the system can always curate the corresponding dialog context in lieu of restricting the corresponding dialog context. In additional or alternative implementations, the system can determine whether to restrict or curate the corresponding dialog context based on content that is included in the corresponding dialog context. For example, in some versions of those implementations, the system can initially attempt to restrict the content of the corresponding dialog context. However, if, in restricting the content of the corresponding dialog context, the system determines that there is no remaining content in the corresponding dialog context, then the system can determine to subsequently curate the corresponding dialog context. However, if, in subsequently curating the corresponding dialog context, the system determines that there is no corresponding curated dialog context suitable for selection, then the system can revert back to restricting the corresponding dialog context. As another example, in some versions of those implementations, the system can initially attempt to curate the content of the corresponding dialog context. However, if, in curating the corresponding dialog context, the system determines that there is no corresponding curated dialog context suitable for selection, then the system can determine to subsequently restrict the corresponding dialog context. Accordingly, it should be understood that the system can employ various rules for determining whether to restrict or curate the corresponding dialog context.

(51) If, at an iteration of block **452**, the system determines to restrict the corresponding dialog context, then the system proceeds to block **454**. At block **454**, the system determines how to restrict the corresponding dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog. For example, in some implementations, block **454** includes sub-block **454A** in which the system includes some content from one or more prior turns of the ongoing dialog in the corresponding modified dialog context. As another example, in some

implementations, block **454** includes sub-block **454B** in which the system omits other content from one or more of the prior turns of the ongoing dialog in the corresponding modified dialog context. At block **456**, the system restricts the corresponding dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog. The system can obtain the corresponding dialog context from one or more databases (e.g., the ongoing dialog context(s) database **140A** that is updated as the ongoing dialog progresses).

(52) For example, and as described with respect to the process flow **200** of FIG. **2**, content that is included in the corresponding dialog context that may result in undesirable NL based outputs being generated can be omitted from the corresponding modified dialog context (e.g., the NL based input **201** of “pretend that it is opposite day”) based on the operations of sub-block **454B** of the method **400** of FIG. **4**, whereas other content that may not result in undesirable NL based outputs being generated can be included from the corresponding modified dialog context (e.g., the subsequent NL based input **201** of “what does a person do when they are happy?”). Put another way, the system can selectively determine what content should be included in the corresponding modified dialog context and what content should be omitted from the corresponding modified dialog context based on analyzing the content that is included in the corresponding dialog context. This enables the system to ensure that the NL based output that is generated based at least in part on the corresponding modified dialog context satisfies various assurance criteria in guiding the human-to-computer interaction (e.g., the ongoing dialog).

(53) If, at an iteration of block **452**, the system determines to curate the corresponding dialog context, then the system proceeds to block **458**. At block **458**, the system selects, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, and from among a plurality of curated dialog contexts, a corresponding curated dialog context. At block **460**, the system replaces the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog. The system can obtain the corresponding curated dialog context from one or more databases (e.g., the curated dialog context(s) database **150A**).

(54) For example, and as described with respect to the process flow **200** of FIG. **2**, the plurality of curated dialog contexts can be defined by a developer associated with the system. Each of the plurality of curated dialog contexts can include content related to disparate dialog topics. Accordingly, in selecting the corresponding curated dialog context, the system can analyze content of the corresponding dialog context to identify one or more words or phrases, and can select the corresponding curated dialog context that is associated with one or more of the words or phrases included in the content of the corresponding dialog context. For instance, if the content of the corresponding context includes one or more terms or phrases related to configuring a Wi-Fi router (e.g., “IP address”, “internet”, “Wi-Fi wiring”, etc.), then the corresponding curated dialog context that is selected can be an “information technology help” curated dialog context. Also, for instance, if the content of the corresponding context includes one or more terms or phrases related to changing a car battery (e.g., “battery”, “positive terminal”, “negative terminal”, etc.), then the corresponding curated dialog context that is selected can be an “car maintenance help” curated dialog context. Further, the system can replace the content of the corresponding dialog context with the content of the selected corresponding curated dialog context to generate the corresponding modified dialog context. Similarly, this enables the system to ensure that the NL based output that is generated based at least in part on the corresponding modified dialog context satisfies various assurance criteria in guiding the human-to-computer interaction (e.g., the ongoing dialog).

(55) Although the method **400** of FIG. **4** depicts particular techniques for managing the LLM through modifying the corresponding dialog context, it should be understood that is for the sake of example and is not meant to be limiting. For example, as described with respect to FIG. **5**, other techniques are contemplated herein for managing the LLM. These other techniques also ensure that

the NL based output that is generated based at least in part on the corresponding dialog context satisfies various assurance criteria in guiding the human-to-computer interaction (e.g., the ongoing dialog), but not require that the system actually modify the corresponding dialog context.

(56) Turning now to FIG. 5, a flowchart illustrating an example method **500** of modifying a corresponding natural language (NL) based output threshold that is utilized in determining NL based output and based on large language model (LLM) output generated using an LLM during an ongoing dialog is depicted. For convenience, the operations of the method **500** are described with reference to a system that performs the operations. This system of the method **500** includes one or more processors, memory, and/or other component(s) of computing device(s) (e.g., client device **110** of FIG. 1, NL based output system **120** of FIG. 1, computing device **610** of FIG. 6, one or more servers, and/or other computing devices). Moreover, while operations of the method **500** are shown in a particular order, this is not meant to be limiting. One or more operations may be reordered, omitted, and/or added.

(57) At block **552**, the system receives NL based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog. At block **554**, the system generates, based on processing the NL based input using an LLM, NL based output that is responsive to the NL based input. At block **556**, the system causes the NL based output to be rendered at the client device. At block **558**, the system receives subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog. The operations of block **552-558** can be performed in the same or similar manner described with respect to the operations of block **352-358**, respectively, of the method **300** of FIG. 3.

(58) At block **560**, the system determines whether to modify a corresponding NL based output threshold. The system can determine whether to modify the corresponding NL based output threshold in the same or similar manner described with respect to the system determining whether to modify the corresponding dialog context as described with respect to the operations of block **360** of the method **300** of FIG. 3. In various implementations, whether the system determines whether to modify the corresponding dialog context as described with respect to the operations of block **360** of the method **300** of FIG. 3 or determines whether to modify the corresponding NL based output threshold can be based on how the system is configured. However, it should be noted that any modified NL based output threshold can be stored in association with a user of the client device from which the NL based input was received at the operations of block **552** of the method **500** of FIG. 5 (e.g., in association with a user identifier for the user of the client device). This enables the system to utilize the modified NL based output threshold in subsequent dialogs that are initiated by the user.

(59) If, at an iteration of block **560**, the system determines to modify the corresponding NL based output threshold, then the system proceeds to block **562**. At block **562**, the system modifies a corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating subsequent NL based output that is responsive to the subsequent NL based input. At block **564**, the system generates, based on processing the subsequent NL based input using the LLM and based on the corresponding modified NL based output threshold that is responsive to the subsequent NL based input. The corresponding NL based output threshold can be, for example, any of the one or more ranking criteria utilized in selecting the words or phrases for inclusion in the NL based output (e.g., as described with respect to the operations of block **354** of the method **300** of FIG. 3), such as an assurance criterion, an accuracy criterion, a quality criterion, and/or any other ranking criterion. By modifying one or more thresholds associated with one or more of these ranking criteria, the subsequent NL based output that is generated based on processing the subsequent NL based input and using the LLM reflects higher assurance, higher accuracy, higher quality etc. than would otherwise be permissible by the system.

(60) For instance, in implementations where the assurance criterion is increased, the system can

ensure that the level of assurance or safety associated with each of the word or phrases selected for inclusion in the subsequent NL based output reflect a higher level of assurance or safety. Further, in additional or alternative implementations where the accuracy criterion is increased, the system can ensure that the level of accuracy or trustworthiness associated with each of the words or phrases reflect a higher level of accuracy. Moreover, in additional or alternative implementations where the accuracy criterion is increased, the system can ensure that the level of quality associated with each of the words or phrases reflect a higher level of quality.

(61) In some implementations, the corresponding NL based output threshold can be increased by a fixed amount (e.g., the words or phrases selected for inclusion in the subsequent NL based output must satisfy a 0.9 threshold assurance threshold rather than a 0.7 assurance threshold). In other implementations, the corresponding NL based output threshold can be increased by a dynamic amount (e.g., the words or phrases selected for inclusion in the subsequent NL based output must satisfy a 0.9 threshold assurance threshold in some instances, but a 0.8 assurance threshold in other instances). In these implementations, the system can determine the dynamic amount to adjust the corresponding NL based output threshold based on, for example, content of a corresponding dialog context for the ongoing dialog. For instance, in implementations where the content of the corresponding dialog context indicate that harm to person or property may result from a user acting upon the NL based output threshold, the corresponding NL based output threshold can be increased by a larger amount.

(62) At block **566**, the system causes the subsequent NL based output to be rendered at the client device. The system can cause the subsequent NL based output to be rendered at the client device in the same or similar manner described with respect to the operations of block **356** of the method **300** of FIG. **3**. Further, the system can return to block **558** and continue processing any further subsequent NL based inputs in the same or similar manner described herein. However, it should be noted that the corresponding dialog continue will continue to be built as the ongoing dialog progresses. Accordingly, the system may perform additional analysis of the corresponding dialog context as it continues to be built as the ongoing dialog progresses (e.g., as described with respect to the process flow **200** of FIG. **2**). In some implementations, the system can continue to utilize the modified NL based output threshold as the ongoing dialog progresses. In these implementations, the system can omit the operations of block **560** of FIG. **5** for each of the subsequent iterations of the method. However, in additional or alternative implementations, the system can perform the operations of block **560** of FIG. **5** for each of the subsequent iterations of the method.

(63) If, at an iteration of block **560**, the system determines not to modify the corresponding NL based output threshold, then the system proceeds to block **568**. At block **568**, the system generates, based on processing the subsequent NL based input using the LLM and based on the corresponding NL based output threshold, subsequent NL based output that is responsive to the subsequent NL based input. The system can generate the subsequent NL based output that is responsive to the subsequent NL based input in the same or similar manner described above with respect to the operations of block **354** of the method **300** of FIG. **3**. However, in generating the subsequent NL based output according to the operations of block **568**, the system may utilize an unmodified NL based output threshold. Notably, in generating the subsequent NL based output according to the operations of block **568**, and in contrast with generating the subsequent NL based output according to the operations of block **564**, the system generates the subsequent NL based output with less stringent NL based output threshold. The difference between the more stringent and less stringent NL based output thresholds will result in the subsequent NL based output that is generated according to the operations of block **568** to differ from the subsequent NL based output that is generated according to the operations of block **564**. This difference is due at least in part to the difference in the corresponding NL based output thresholds.

(64) At block **570**, the system causes the subsequent NL based output to be rendered at the client device. The system can cause the subsequent NL based output to be rendered at the client device in

the same or similar manner described with respect to the operations of block **356** of the method **300** of FIG. **3**. Further, the system can return to block **558** and continue processing any further subsequent NL based inputs in the same or similar manner described herein. However, it should be noted that the corresponding dialog continue will continue to be built as the ongoing dialog progresses. Accordingly, the system may perform additional analysis of the corresponding dialog context as it continues to be built as the ongoing dialog progresses (e.g., as described with respect to the process flow **200** of FIG. **2**). In some implementations, the system can continue to utilize the modified NL based output threshold as the ongoing dialog progresses. In these implementations, the system can omit the operations of block **560** of FIG. **5** for each of the subsequent iterations of the method. However, in additional or alternative implementations, the system can perform the operations of block **560** of FIG. **5** for each of the subsequent iterations of the method.

(65) Although the method **300** of FIG. **3** and the method **500** of FIG. **5** depict different techniques for managing the LLM, it should be understood that is for the sake of illustrating different techniques described herein and is not meant to be limiting. For example, these techniques can optionally be utilized in combination with one another.

(66) Turning now to FIG. **6**, a block diagram of an example computing device **610** that may optionally be utilized to perform one or more aspects of techniques described herein is depicted. In some implementations, one or more of a client device, cloud-based automated assistant component(s) or other cloud-based software application component(s), and/or other component(s) may comprise one or more components of the example computing device **610**.

(67) Computing device **610** typically includes at least one processor **614** which communicates with a number of peripheral devices via bus subsystem **612**. These peripheral devices may include a storage subsystem **624**, including, for example, a memory subsystem **625** and a file storage subsystem **626**, user interface output devices **620**, user interface input devices **622**, and a network interface subsystem **616**. The input and output devices allow user interaction with computing device **610**. Network interface subsystem **616** provides an interface to outside networks and is coupled to corresponding interface devices in other computing devices.

(68) User interface input devices **622** may include a keyboard, pointing devices such as a mouse, trackball, touchpad, or graphics tablet, a scanner, a touch screen incorporated into the display, audio input devices such as voice recognition systems, microphones, and/or other types of input devices. In general, use of the term “input device” is intended to include all possible types of devices and ways to input information into computing device **610** or onto a communication network.

(69) User interface output devices **620** may include a display subsystem, a printer, a fax machine, or non-visual displays such as audio output devices. The display subsystem may include a cathode ray tube (CRT), a flat-panel device such as a liquid crystal display (LCD), a projection device, or some other mechanism for creating a visible image. The display subsystem may also provide non-visual display such as via audio output devices. In general, use of the term “output device” is intended to include all possible types of devices and ways to output information from computing device **610** to the user or to another machine or computing device.

(70) Storage subsystem **624** stores programming and data constructs that provide the functionality of some or all of the modules described herein. For example, the storage subsystem **624** may include the logic to perform selected aspects of the methods disclosed herein, as well as to implement various components depicted in FIGS. **1** and **2**.

(71) These software modules are generally executed by processor **614** alone or in combination with other processors. Memory **625** used in the storage subsystem **624** can include a number of memories including a main random access memory (RAM) **630** for storage of instructions and data during program execution and a read only memory (ROM) **632** in which fixed instructions are stored. A file storage subsystem **626** can provide persistent storage for program and data files, and may include a hard disk drive, a floppy disk drive along with associated removable media, a CD-ROM drive, an optical drive, or removable media cartridges. The modules implementing the

functionality of certain implementations may be stored by file storage subsystem **626** in the storage subsystem **624**, or in other machines accessible by the processor(s) **614**.

(72) Bus subsystem **612** provides a mechanism for letting the various components and subsystems of computing device **610** communicate with each other as intended. Although bus subsystem **612** is shown schematically as a single bus, alternative implementations of the bus subsystem **612** may use multiple busses.

(73) Computing device **610** can be of varying types including a workstation, server, computing cluster, blade server, server farm, or any other data processing system or computing device. Due to the ever-changing nature of computers and networks, the description of computing device **610** depicted in FIG. **6** is intended only as a specific example for purposes of illustrating some implementations. Many other configurations of computing device **610** are possible having more or fewer components than the computing device depicted in FIG. **6**.

(74) In situations in which the systems described herein collect or otherwise monitor personal information about users, or may make use of personal and/or monitored information), the users may be provided with an opportunity to control whether programs or features collect user information (e.g., information about a user's social network, social actions or activities, profession, a user's preferences, or a user's current geographic location), or to control whether and/or how to receive content from the content server that may be more relevant to the user. Also, certain data may be treated in one or more ways before it is stored or used, so that personal identifiable information is removed. For example, a user's identity may be treated so that no personal identifiable information can be determined for the user, or a user's geographic location may be generalized where geographic location information is obtained (such as to a city, ZIP code, or state level), so that a particular geographic location of a user cannot be determined. Thus, the user may have control over how information is collected about the user and/or used.

(75) In some implementations, a method implemented by one or more processors is provided, and includes receiving natural language (NL) based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog; generating, based on processing the NL based input using a large language model (LLM), NL based output that is responsive to the NL based input; causing the NL based output that is responsive to the NL based input to be rendered at the client device; receiving subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog; determining, based on at least the NL based output and/or the subsequent NL based input, whether to modify a corresponding dialog context for the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding context for the given subsequent turn of the ongoing dialog: modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog; generating, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM, subsequent NL based output that is responsive to the subsequent NL based input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

(76) These and other implementations of technology disclosed herein can optionally include one or more of the following features.

(77) In some implementations, modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: restricting the corresponding dialog context to a given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(78) In some versions of those implementations, restricting the corresponding dialog context to the

given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: omitting at least the given turn of the ongoing dialog from the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(79) In some implementations, modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: supplanting the corresponding dialog context with a corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(80) In some versions of those implementations, supplanting the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: selecting, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, and from among a plurality of curated dialog contexts, the corresponding curated dialog context.

(81) In some implementations, determining whether to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog based on the NL based output and/or the subsequent NL based input can include: determining, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, an assurance score for the given subsequent turn of the ongoing dialog; and determining, based on the assurance score, whether to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog.

(82) In some versions of those implementations, determining to modify the corresponding context for the given subsequent turn of the ongoing dialog can be in response to determining that the assurance score fails to satisfy an assurance score threshold.

(83) In additional or alternative versions of those implementations, determining the assurance score for the given subsequent turn of the ongoing dialog based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input can include: processing, using an assurance machine learning (ML) model, the corresponding output content captured in the NL based output and/or the corresponding subsequent input content captured in the subsequent NL based input to generate output; and determining, based on the output generated using the assurance ML model, the assurance score.

(84) In some implementations, the method can further include determining the corresponding dialog context for the given subsequent turn of the ongoing dialog based on at least the NL based input, the NL based output that is responsive to the NL based input, and the subsequent NL based input.

(85) In some versions of those implementations, determining the corresponding dialog context for the given subsequent turn of the ongoing dialog can be further based on a dialog history of a user of the client device.

(86) In some implementations, the method can further include: determining a prior corresponding dialog context for the given turn of the ongoing dialog based on at least the NL based input; and determining, based on at least the NL based input, whether to modify the prior corresponding dialog context for the given turn of the ongoing dialog.

(87) In some versions of those implementations, determining the prior corresponding dialog context for the given turn of the ongoing dialog can be further based on a dialog history of a user of the client device.

(88) In some implementations, the method can further include: determining, based on the NL based output and/or the subsequent NL based input, whether to modify a corresponding NL based output threshold for generating subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; and in response to

determining to modify the corresponding NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; modifying the corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.

(89) In some versions of those implementations, generating the subsequent NL based output can be further based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.

(90) In additional or alternative versions of those implementations, the method can further include storing the corresponding modified NL based output threshold in a dialog history of a user of the client device; and utilizing the corresponding modified NL based output threshold in conducting future dialogs with the user of the client device.

(91) In some implementations, the method can further include, in response to determining to refrain from modifying the corresponding context for the given subsequent turn of the ongoing dialog: generating, based on processing the subsequent NL based input and the corresponding dialog context for the given subsequent turn of the ongoing dialog using the LLM, the subsequent NL based output that is responsive to the subsequent NL based input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

(92) In some versions of those implementations, the subsequent NL based output generated based on processing the subsequent NL based input and the corresponding dialog context for the given subsequent turn of the ongoing dialog using the LLM can differ from the subsequent NL based output generated based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM due to a difference between the corresponding dialog context and the corresponding modified dialog context.

(93) In some implementations, a method implemented by one or more processors is provided, and includes receiving natural language (NL) based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog; generating, based on processing the NL based input using a large language model (LLM), NL based output that is responsive to the NL based input; causing the NL based output that is responsive to the NL based input to be rendered at the client device; receiving subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog; determining, based on the NL based output and/or the subsequent NL based input, whether to modify a corresponding NL based output threshold for generating subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog: modifying the corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; generating, based on processing the subsequent NL based input using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

(94) These and other implementations of technology disclosed herein can optionally include one or

more of the following features.

(95) In some implementations, modifying the corresponding NL based output threshold to generate the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog can include increasing an assurance score threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.

(96) In some implementations, the method can further include storing the corresponding modified NL based output threshold in a dialog history of a user of the client device; and utilizing the corresponding modified NL based output threshold in conducting future dialogs with the user of the client device.

(97) In some implementations, the method can further include determining, based on at least the NL based output and/or the subsequent NL based input, whether to modify a corresponding dialog context for the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding context for the given subsequent turn of the ongoing dialog: modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog; generating, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

(98) In some versions of those implementations, modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include restricting the corresponding dialog context to a given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(99) In some further versions of those implementations, restricting the corresponding dialog context to the given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: omitting at least the given turn of the ongoing dialog from the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(100) In additional or alternative versions of those implementations, modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include: supplanting the corresponding dialog context with a corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

(101) In some additional or alternative versions of those implementations, supplanting the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog can include selecting, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, and from among a plurality of curated dialog contexts, the corresponding curated dialog context.

(102) In addition, some implementations include one or more processors (e.g., central processing unit(s) (CPU(s)), graphics processing unit(s) (GPU(s), and/or tensor processing unit(s) (TPU(s)) of one or more computing devices, where the one or more processors are operable to execute

instructions stored in associated memory, and where the instructions are configured to cause performance of any of the aforementioned methods. Some implementations also include one or more computer readable storage media (e.g., transitory and/or non-transitory) storing computer instructions executable by one or more processors to perform any of the aforementioned methods. Some implementations also include a computer program product including instructions executable by one or more processors to perform any of the aforementioned methods.

Claims

1. A method implemented by one or more processors, the method comprising: receiving natural language (NL) based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog; generating, based on processing the NL based input using a large language model (LLM), NL based output that is responsive to the NL based input; causing the NL based output that is responsive to the NL based input to be rendered at the client device; receiving subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog; determining, based on the NL based output and/or the subsequent NL based input, whether to modify a corresponding NL based output threshold for generating subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog: modifying the corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; generating, based on processing the subsequent NL based input using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.
2. The method of claim 1, wherein modifying the corresponding NL based output threshold to generate the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog comprises: increasing an assurance score threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.
3. The method of claim 2, further comprising: storing the corresponding modified NL based output threshold in a dialog history of a user of the client device; and utilizing the corresponding modified NL based output threshold in conducting future dialogs with the user of the client device.
4. The method of claim 1, further comprising: determining, based on at least the NL based output and/or the subsequent NL based input, whether to modify a corresponding dialog context for the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog: modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog; generating, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based

input; and causing the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

5. The method of claim 4, wherein modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprises: restricting the corresponding dialog context to a given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

6. The method of claim 5, wherein restricting the corresponding dialog context to the given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprises: omitting at least the given turn of the ongoing dialog from the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

7. The method of claim 4, wherein modifying the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprises: supplanting the corresponding dialog context with a corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

8. The method of claim 7, wherein supplanting the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprises: selecting, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, and from among a plurality of curated dialog contexts, the corresponding curated dialog context.

9. A system comprising: at least one processor; and memory storing instructions that, when executed, cause the at least one processor to be operable to: receive natural language (NL) based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog; generate, based on processing the NL based input using a large language model (LLM), NL based output that is responsive to the NL based input; cause the NL based output that is responsive to the NL based input to be rendered at the client device; receive subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog; determine, based on the NL based output and/or the subsequent NL based input, whether to modify a corresponding NL based output threshold for generating subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog: modify the corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; generate, based on processing the subsequent NL based input using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and cause the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

10. The system of claim 9, wherein the instructions to modify the corresponding NL based output threshold to generate the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog comprise instructions to: increase an assurance

score threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.

11. The system of claim 10, wherein the at least one processor is further operable to: store the corresponding modified NL based output threshold in a dialog history of a user of the client device; and utilize the corresponding modified NL based output threshold in conducting future dialogs with the user of the client device.

12. The system of claim 9, wherein the at least one processor is further operable to: determine, based on at least the NL based output and/or the subsequent NL based input, whether to modify a corresponding dialog context for the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog: modify the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog; generate, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and cause the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

13. The system of claim 12, wherein the instructions to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprise instructions to: restrict the corresponding dialog context to a given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

14. The system of claim 13, wherein the instructions to restrict the corresponding dialog context to the given prior turn of the ongoing dialog that occurred prior to the given turn of an ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprise instructions to: omit at least the given turn of the ongoing dialog from the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

15. The system of claim 12, wherein the instructions to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprise instructions to: supplant the corresponding dialog context with a corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog.

16. The system of claim 15, wherein the instructions to supplant the corresponding dialog context with the corresponding curated dialog context to generate the corresponding modified dialog context for the given subsequent turn of the ongoing dialog comprise instructions to: select, based on corresponding output content captured in the NL based output and/or corresponding subsequent input content captured in the subsequent NL based input, and from among a plurality of curated dialog contexts, the corresponding curated dialog context.

17. A non-transitory computer-readable storage medium storing instructions that, when executed by at least one processor, cause the at least one processor to perform operations to: receive natural language (NL) based input associated with a client device, the NL based input being received during a given turn of an ongoing dialog; generate, based on processing the NL based input using a large language model (LLM), NL based output that is responsive to the NL based input; cause the NL based output that is responsive to the NL based input to be rendered at the client device; receive subsequent NL based input associated with the client device, the subsequent NL based input being received during a given subsequent turn of the ongoing dialog; determine, based on the NL based output and/or the subsequent NL based input, whether to modify a corresponding NL based output

threshold for generating subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog: modify the corresponding NL based output threshold to generate a corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog; generate, based on processing the subsequent NL based input using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and cause the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.

18. The non-transitory computer-readable storage medium of claim 17, wherein the operations to modify the corresponding NL based output threshold to generate the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog comprise operations to: increase an assurance score threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog.

19. The non-transitory computer-readable storage medium of claim 18, wherein instructions further cause the at least one processor to perform operations to store the corresponding modified NL based output threshold in a dialog history of a user of the client device; and utilize the corresponding modified NL based output threshold in conducting future dialogs with the user of the client device.

20. The non-transitory computer-readable storage medium of claim 17, wherein instructions further cause the at least one processor to perform operations to determine, based on at least the NL based output and/or the subsequent NL based input, whether to modify a corresponding dialog context for the given subsequent turn of the ongoing dialog; and in response to determining to modify the corresponding dialog context for the given subsequent turn of the ongoing dialog: modify the corresponding dialog context for the given subsequent turn of the ongoing dialog to generate a corresponding modified dialog context for the given subsequent turn of the ongoing dialog; generate, based on processing the subsequent NL based input and the corresponding modified dialog context for the given subsequent turn of the ongoing dialog using the LLM and based on the corresponding modified NL based output threshold for generating the subsequent NL based output that is responsive to the subsequent NL based input received during the given subsequent turn of the ongoing dialog, the subsequent NL based output that is responsive to the subsequent NL based input; and cause the subsequent NL based output that is responsive to the subsequent NL based input to be rendered at the client device.
