

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12406176
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Kolhe; Jitendra Onkar et al.

Distributed execution of neural networks

Abstract

Example techniques for causing execution of neural networks are described. A neural network includes a first part and a second part. A determination is made that a first physical resource in a first computing device is to execute the first part and that a second physical resource in a second computing device is to execute the second part. The determination is based on a latency in communication between the first physical resource and the second physical resource. The first computing device and the second computing device are part of a cluster.

Inventors: Kolhe; Jitendra Onkar (Karnataka, IN), Chatterjee; Soumitra (Karnataka, IN), Parthasarathy; Mohan (Karnataka, IN)

Applicant: Hewlett Packard Enterprise Development LP (Houston, TX)

Family ID: 82024063

Assignee: Hewlett Packard Enterprise Development LP (Spring, TX)

Appl. No.: 17/231153

Filed: April 15, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220198249 A1	Jun. 23, 2022

Foreign Application Priority Data

IN	2020041055302	Dec. 18, 2020
----	---------------	---------------

Publication Classification

Int. Cl.: G06N3/063 (20230101); G06N3/10 (20060101)

U.S. Cl.:

CPC **G06N3/063** (20130101); **G06N3/10** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7426578	12/2007	Jones	709/248	H04L 67/1004
10504022	12/2018	Khaitan et al.	N/A	N/A
11182314	12/2020	Borkovic	N/A	G06N 3/063
11216314	12/2021	Harwood	N/A	G06F 9/4856
11452122	12/2021	Macias	N/A	H04W 72/1263
2017/0273005	12/2016	Daragon	N/A	H04W 40/246
2019/0050717	12/2018	Temam	N/A	G06N 3/048
2019/0095796	12/2018	Chen	N/A	G06N 3/045
2019/0114535	12/2018	Ng et al.	N/A	N/A
2019/0188570	12/2018	Aldea Lopez	N/A	G06N 3/084
2020/0202215	12/2019	Sadowski	N/A	G06N 3/063
2021/0295139	12/2020	Kolhe	N/A	G06N 3/045
2022/0052961	12/2021	Chauhan	N/A	H04L 47/82
2022/0116452	12/2021	Rafey	N/A	G06F 18/24137

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
3367310	12/2017	EP	N/A

OTHER PUBLICATIONS

Ankit et al., "PUMA: A Programmable Ultra-efficient Memristor-based Accelerator for Machine Learning Inference", ACM, Apr. 13-17, 2019, 17 pages. cited by applicant

Avnet, "AI at the Edge: The next frontier of the Internet of Things", 2018, 10 pages. cited by applicant

HPE, "HPE Converged Edge Systems Video Analytics and Management Solution", Part No. P09283-002, Jun. 2019, 31 pages. cited by applicant

Qiao et al., "DLTAP: A Network-efficient Scheduling Method for Distributed Deep Learning Workload in Containerized Cluster Environment", ITM Web of Conferences 12, 03030, Sep. 5, 2017, 5 Pages. cited by applicant

Stack Exchange, "Shortest path with exactly k edges", available online at <<https://cs.stackexchange.com/questions/11503/shortest-path-with-exactly-k-edges>>, 2021, 2 pages. cited by applicant

Tomkins et al., "Google AI Blog", available online at <<https://web.archive.org/web/20200802122231/https://ai.googleblog.com/2017/11/>>, Nov. 30, 2017, 38 pages. cited by applicant

Zagoruyko et al., "Wide Residual Networks", 2016, 15 pages. cited by applicant

Background/Summary

BACKGROUND

(1) Artificial neural networks, also referred to as neural networks, are simplified models of the way a human nervous system operates. A neural network may include several interconnected nodes arranged in a plurality of layers. Neural networks may be used to implement, for example, pattern recognition and predictive analytics in a variety of applications, such as speech recognition, machine translation, computer vision, and healthcare diagnostics.

Description

BRIEF DESCRIPTION OF DRAWINGS

- (1) The following detailed description references the figures, wherein:
- (2) FIG. 1 illustrates a computing device for causing execution of neural networks, according to an example implementation of the present subject matter;
- (3) FIG. 2 illustrates a computing environment that includes a plurality of computing devices that are to execute a neural network, according to an example implementation of the present subject matter;
- (4) FIG. 3 illustrates a computing environment implementing scheduling of parts of neural network for execution of a neural network, according to an example implementation of the present subject matter;
- (5) FIG. 4 illustrates the connections among physical resources of a first computing device and a second computing device, according to an example implementation of the present subject matter;
- (6) FIG. 5 illustrates a cluster latency table providing information of latency between physical resources of computing devices, according to an example implementation of the present subject matter;
- (7) FIG. 6 illustrates a cluster having a first computing device and a second computing device, according to an example implementation of the present subject matter;
- (8) FIG. 7 illustrates a method for causing execution of neural networks, according to an example implementation of the present subject matter;
- (9) FIG. 8 illustrates a method for scheduling parts of a neural network to physical resources, according to an example implementation of the present subject matter; and
- (10) FIG. 9 illustrates a computing environment implementing a non-transitory computer-readable medium for causing execution of neural networks, according to an example implementation of the present subject matter.

DETAILED DESCRIPTION

(11) A neural network may include a plurality of instructions, such as instructions to perform matrix-vector multiplications, instructions to perform vector-vector additions, instructions to perform scalar-scalar additions, instructions to load data, instructions to store data, and the like. The instructions may be executed by a physical resource, such as a central processing unit (CPU), a graphics processing unit (GPU), or an accelerator device. An accelerator may refer to a specialized hardware that is capable of performing a particular function (in this case, execution of neural networks) more efficiently as compared to other types of computing architecture. An example

accelerator is a dot-product engine (DPE), which includes a plurality of memristors and is capable of performing matrix-vector multiplication operations as dot product computations.

(12) In some cases, the execution of a neural network, such as a neural network of a large size, may be a computationally-intensive task, and may not be performed as a whole by a single physical resource. In such cases, generally, the neural network may be executed in parts by the physical resource, causing an increase in the execution time of the neural network. This may be because the execution of the neural network in parts may involve repeated reconfiguration of data of the neural network, such as weights of the neural network, on the physical resource. For instance, initially, data corresponding to a first part is to be configured on the physical resource to allow execution of the first part. Upon execution of the first part, data corresponding to a second part is to be configured on the physical resource, so that the second part is executed. The reconfigurations also cause expenditure of computational resources.

(13) The present subject matter relates to execution of neural networks. With the implementations of the present subject matter, a neural network may be executed on a plurality of physical resources. The physical resources may be distributed across computing devices of a duster.

(14) In accordance with an example implementation, a neural network may include a first part and a second part. The first part and the second part may be executed by different physical resources, hi an example, each part may include a set of instructions. Further, each of part may have associated data, such as weights and biases, on which the set of instructions are to operate. For instance, the first part may include instructions to perform matrix-vector multiplications using a weight that corresponds to the first part or instructions to perform vector-vector additions using a bias corresponding to the first part.

(15) Data from execution of one part may be provided as an input to another part. For instance, the neural network may be a multilayer neural network and each part may include instructions corresponding to one layer in the neural network. The result from execution of one layer of the neural network may be transmitted to the next layer in the neural network. As another example, one part may include instructions to perform a matrix-vector multiplication and another part may include instructions to perform an addition of a vector generated from the matrix-vector multiplication with another vector.

(16) A computing device may determine physical resources that are to execute the parts of the neural network. For example, the computing device may determine that a first physical resource in a first computing device is to execute the first part and that a second physical resource in a second computing device is to execute the second part. Upon the determination, the parts may be executed on the determined physical resources. In an example, the first computing device is same as the computing device that determines the physical resources that are to execute the parts.

(17) The determination of the physical resources that are to execute the parts may be performed based on a latency in communication between the physical resources. For example, consider that the first physical resource, the second physical resource, and a third physical resource are available for execution of parts of a neural network. Consider also that the latency in communication between the first physical resource and the second physical resource is less than that between the first physical resource and the third physical resource and that between the second physical resource and the third physical resource. Accordingly, the first physical resource and the second physical resource may be selected for execution of the first part and the second part, respectively.

(18) The latency in communication between the physical resources may depend on the communication link that connects the physical resources. The determination of the physical resources for execution of the parts based on the latency ensures that the data transfer that happens as part of execution of the neural network happens with minimal delay, thereby increasing the speed of execution of the neural network.

(19) The first computing device and the second computing device may be part of a duster of computing devices, also referred to as a cluster. The computing devices of the cluster may share

their respective physical resources for execution of parts of neural network. Accordingly, different parts of a single neural network may be executed in physical resources belonging to different computing devices.

(20) In an example, each part of the neural network may have a corresponding logical identifier. Further, to cause execution of the neural network by a determined physical resource, the logical identifier in the neural network may be substituted with an identifier of the physical resource. In an example, a physical resource may be a dot-product engine (DPE) and the execution of a part may be performed by a tile of the DPE. Further, the identifier of the physical resource may be a physical address of a tile of the DPE that is to execute the part.

(21) The formation of a cluster increases the number of physical resources that are available to execute a neural network. Therefore, speed of execution of a large-sized neural network, which may have several parts, may be significantly increased. For instance, the number of reconfigurations of data of the neural network, such as weights and biases, to be performed on the physical resources may be reduced, as the different parts may be assigned to different physical resources. The present subject matter can be utilized for execution of neural networks in a computing environment having multiple computing devices, where a particular computing device may not have physical resources sufficient for execution of all parts of a neural network.

(22) The following description refers to the accompanying drawings. Wherever possible, the same reference numbers are used in the drawings and the following description to refer to the same or similar parts. While several examples are described in the description, modifications, adaptations, and other implementations are possible and are intended to be covered herein.

(23) FIG. 1 illustrates a computing device **100** for causing execution of neural networks, according to an example implementation of the present subject matter. The computing device **100** may be implemented as, for example, a desktop computer, a laptop computer, or a server. The computing device **100** includes a processor **102** and a memory **104** coupled to the processor **102**.

(24) The processor **102** may be implemented as a microprocessor, a microcomputer, a microcontroller, a digital signal processor, a central processing unit, a state machine, a logic circuitry, and/or any device that can manipulate signals based on operational instructions. Among other capabilities, the processor **102** may fetch and execute computer-readable instructions included in the memory **104**. The computer-readable instructions may include instructions **106-110**. The functions of the processor **102** may be provided through the use of dedicated hardware as well as hardware capable of executing machine-readable instructions.

(25) The memory **104** may include any non-transitory computer-readable medium including volatile memory (e.g., RAM), and/or non-volatile memory (e.g., EPROM, flash memory, Memristor, etc.). The memory **104** may also be an external memory unit, such as a flash drive, a compact disk drive, an external hard disk drive, or the like.

(26) In addition to the processor **102** and the memory **104**, the computing device **100** may also include interface(s) and system data (not shown in FIG. 1). The interface(s) may include a variety of instructions-based interfaces and hardware interfaces that allow interaction with a user and with other communication and computing devices, such as network entities, web servers, external repositories, and peripheral devices. The system data may serve as a repository for storing data that may be fetched, processed, received, or created by the instructions.

(27) In operation, the instructions **106** are executable to receive a neural network including a first set of instructions and a second set of instructions. The first set of instructions and the second set of instructions may be interchangeably referred to as a first part and a second part respectively. The instructions may include, for example, instructions to perform matrix-vector multiplications, instructions to perform vector-vector additions, and instructions to load, store, send, and/or receive data. In an example, output of execution of the first part may have to be fed as input to the second part. For instance, the first part may include instructions corresponding to one layer of the neural network and the second part may include instructions corresponding to the next layer of the neural

network. As another example, the first part may include instructions to perform a matrix-vector multiplication and the second part may include instructions to perform an addition of a vector generated from the matrix-vector multiplication with another vector. The instructions in the neural network may also be referred to as neural network instructions.

(28) The instructions **108** are executable to determine physical resources that are to execute the parts of the neural network. For example, it may be determined that a first physical resource in a first computing device (not shown in FIG. 1) is to execute the first part and that a second physical resource in a second computing device (not shown in FIG. 1) is to execute the second part. The first computing device and the second computing device may be part of a cluster of computing devices, also referred to as a cluster. Each computing device of the cluster may have one or more physical resources. Further, physical resources from multiple computing devices may be utilized in execution of a neural network.

(29) The determination of the physical resources to execute the parts may be performed based on a latency in communication between the first physical resource and the second physical resource. For instance, of multiple physical resources that are available for execution of parts, the communication between the first physical resource and the second physical resource may have the least latency. The latency in communication between the physical resources may depend on the communication link that connects the physical resources. In an example, to determine the latency in communication between two physical resources, the computing device **100** may utilize cluster latency information (not shown in FIG. 1), which may indicate the latencies in communication between physical resources of the cluster.

(30) Upon the determination, the instructions **110** may cause execution of the first part by the first physical resource and the second part by the second physical resource. For instance, instructions of the first part and of the second part may be transmitted to the first physical resource and the second physical resource, respectively, for being executed. The execution of a neural network may be performed for training of the neural network or for inferencing of the neural network.

(31) FIG. 2 illustrates a computing environment **200** that includes a plurality of computing devices that are to execute a neural network, according to an example implementation of the present subject matter. In an example, the computing devices of the computing environment **200** may be converged infrastructure devices, in which compute, storage, and networking are integrated. Further, in an example, the computing devices may be geographically proximate to each other. For instance, the computing environment **200** may be an industrial plant, a factory, an offshore drilling rig, or the like.

(32) The computing devices may have to perform various functions for which a neural network has to be executed. The functions may be, for example, asset tracking, speech recognition, surveillance, and the like. To execute the neural networks, each computing device may include a physical resource. The physical resource may be an accelerator, which may refer to a hardware component that is specialized in the execution of neural networks or parts thereof. In an example, a physical resource may be a dot-product engine (DPE). A physical resource may be accommodated, for example, in a Peripheral Component Interconnect (PCI) slot of a computing device.

(33) While each computing device may have physical resource(s), the number of physical resources in one computing device may not be sufficient to execute a large-sized neural network, which may have several parts. The number of physical resources in the computing device may be small due to various reasons, such as unavailability of PCI slots in the computing device, cost saving, and power saving. To facilitate execution of large-sized neural networks, a cluster **202** having multiple computing devices of the computing environment **200** may be formed. Accordingly, physical resources of multiple computing devices of the cluster **202** may be utilized to execute the neural network. For example, the cluster **202** may include a first computing device **204**, a second computing device **206**, and a third computing device **208**. Further, the first computing device **204** may include a first physical resource **210**, the second computing device **206** may include a second

physical resource **212** and a third physical resource **214**, and the third computing device **208** may include a fourth physical resource **216** and a fifth physical resource **218**. Thus, parts of the neural network may be executed in one or more of the physical resources **210-218**.

(34) The present subject matter provides techniques for determination of physical resources of the cluster **202** that are to be used for execution of parts of a neural network. The determination of the physical resources to execute the parts of the neural network may be referred to as scheduling the parts of the neural network. In an example, the scheduling may be performed by one computing device of the cluster **202**, such as the first computing device **204**. Accordingly, the first computing device **204** corresponds to the computing device **100**. In an example, the scheduling may be performed by a computing device that receives a command to execute the neural network.

(35) The scheduling may be performed by a first scheduler **220** that is part of the first computing device **204**. In an example, the first scheduler **220** may be implemented by a processor of the first computing device **204** by executing a set of instructions (not shown in FIG. 2). Similar to the first scheduler **220**, a second scheduler **222** and a third scheduler **224** may be implemented in the second computing device **206** and the third computing device **208**, respectively.

(36) To perform the scheduling, the first scheduler **220** may receive information regarding the neural network, such as number of parts in the neural network and the sequence in which the various parts are to be executed, and may schedule the different parts for execution. For example, the first scheduler **220** may assign the first part for execution by the first physical resource **210** and the second part for execution by the second physical resource **212**. The scheduling may be based on latencies in communications between the various physical resources, as will be explained below.

(37) Although not shown, the various computing devices of the cluster **202** may communicate with each other over a communication network, which may be a wireless or a wired network, or a combination thereof. The communication network may be a collection of individual networks, interconnected with each other and functioning as a single large network (e.g., the internet or an intranet). Examples of such individual networks include Global System for Mobile Communication (GSM) network, Universal Mobile Telecommunications System (UMTS) network, Personal Communications Service (PCS) network, Time Division Multiple Access (TDMA) network, Code Division Multiple Access (CDMA) network, Next Generation Network (NGN), Public Switched Telephone Network (PSTN), and Integrated Services Digital Network (ISDN). Depending on the technology, the communication network may include various network entities, such as transceivers, gateways, and routers.

(38) FIG. 3 illustrates a computing environment **300** implementing scheduling of parts of neural network for execution of the neural network, according to an example implementation of the present subject matter. The computing environment **300** may correspond to the computing environment **200**. The computing environment **300** may include the first computing device **204** and the second computing device **206**, which form part of the cluster **202** (not shown in FIG. 3).

(39) The first computing device **204** may include a compiler **302** that may compile a neural network model **304** to generate a neural network **306**. The neural network model **304** may be programmed in a high-level programming language, such as C++ or python. The neural network model **304** may be compiled so that the generated neural network **306** is in a machine-readable form and can be executed by physical resources. In an example, the compiler **302** may be implemented by a processor of the first computing device **204** by executing a set of instructions (not shown in FIG. 3).

(40) In an example, the neural network **306** may include parameters, also referred to as neural network parameters, which include weights and biases. A weight of a neural network may be assigned to a connection between a first node and a second node of the neural network, where each node represents an artificial neuron of the neural network. The weight represents a strength of the connection between the nodes. For instance, an output of the first node is multiplied by the weight of the connection between the first node and the second node, and the product is used for obtaining

the output of the second node. In an example, the weight may be represented in the form of a matrix and the output of the first node may be represented in the form of a vector. Accordingly, the multiplication of the weight and the output of the first node may be performed through matrix-vector multiplication (MVM). A bias may refer to an additional parameter that is added to the product of the weight and the output of the first node. The bias may be used to adjust the value output by the second node. In an example, the bias may be represented in the form of a vector. Accordingly, the sum of the product of the MVM and the bias may be a vector-vector addition.

(41) The neural network **306** may also include a plurality of parts, such as a first part **308** (also referred to as a first set of instructions **308**) and a second part **310** (also referred to as a second set of instructions **310**). The sets of instructions **308** and **310** are to be executed to execute the neural network **306**. For example, a set of instructions may include instructions to perform an MVM operation which is to be performed to compute a product of a weight matrix and an output vector from a node, and instructions to perform vector-vector addition which is to be performed to compute a sum of the product of the MVM and a bias. A set of instructions may also include instructions to perform scalar-scalar addition, instructions to perform scalar-scalar multiplication, and instructions to load data, store data, send data, and receive data.

(42) A part of the neural network **306** may be associated with a set of weights, a set of biases, or both. Further, the instructions of the part may operate on the corresponding weights and/or biases. For instance, the first set of instructions **308** may be associated with a first set of weights **316** and a first set of biases **318**. Accordingly, the first set of instructions **308** may perform MVMs involving the first set of weights **316** and vector-vector additions involving the first set of biases **318**.

Similarly, the second set of instructions **310** may be associated with a second set of weights **320** and a second set of biases **322**, and the second set of instructions **310** may perform MVMs involving the second set of weights **320** and vector-vector additions involving the second set of biases **322**. In an example, the second set of instructions **310** may receive an input generated based on execution of the first set of instructions **308**. For instance, an output generated from a first MVM performed as a result of execution of the first set of instructions **308** may act as an input for performing a second MVM that is to be performed by executing the second set of instructions **310**.

(43) A set of instructions and the corresponding weights and biases may be collectively referred to as compute block. For instance, a first compute block **324** may include the first set of instructions **308**, the first set of weights **316**, and the first set of biases **318** and a second compute block **326** may include the second set of instructions **310**, the second set of weights **320**, and the second set of biases **322**.

(44) In an example, the compiler **302** may generate the neural network **306** such that each part of the neural network **306** may be executed by one of the physical resources at any given point in time. Further, the neural network **306** may include a logical identifier corresponding to each part, such as a first logical identifier **328** corresponding to the first part **308** and a second logical identifier **330** corresponding to the second part **310**, to indicate that the corresponding part is to be assigned to one physical resource for execution. The logical identifier may be replaced with the identifier of the physical resource that is to execute the corresponding part. The identifier of the physical resource may be, for example, a physical address of the physical resource. Further, in an example, the generation of the neural network **306** may be such that each part may be executed by a computational unit belonging to a physical resource. The computational unit may be, for example, a the of a DPE. Further, the logical identifier may indicate that the corresponding part is to be assigned to one computational unit for execution. The logical identifier may be replaced with the identifier of the computational unit that is to execute the corresponding part. Thus, the logical identifier may act as a placeholder for the identifier of a physical resource or the identifier of the computational unit, as the case may be.

(45) The scheduling of the parts will be explained. In an example, the scheduling of the parts of a neural network may be performed by the scheduler of the computing device that received an

instruction to execute the neural network, such as from a user. For instance, if an instruction to execute the neural network **306** is received by the first computing device **204**, the scheduling of the parts of the neural network **306** may be performed by the first scheduler **220**. To perform the scheduling, first, the logical identifiers in the neural network **306** may be extracted. The extraction may be performed by a first loader **332** of the first computing device **204**. The first loader **332** may perform the extraction, for example, by parsing the neural network **306**. The first loader **332** may be implemented by the processor of the first computing device **204** by executing a set of instructions (not shown in FIG. 3). The first loader **332** may provide the extracted logical identifiers to the first scheduler **220**.

(46) Upon receiving the logical identifiers, the first scheduler **220** may determine the physical resources to which the parts corresponding to the logical identifiers are to be assigned for execution. The first scheduler **220** may determine the physical resources based on latencies in communication between the physical resources in the cluster **202**. The latencies may be considered for the determination because the execution of the neural network **306** involves flow of data (e.g., output of execution of a part of the neural network **306**) between the various parts of the neural network **306**. Accordingly, data flows between physical resources when executing the parts. The determination of the physical resources based on the latencies ensures that the delay in communication of data between physical resources is minimized, thereby improving the speed of execution of the neural network. To determine the latency in communication between the various physical resources of the cluster **202**, cluster latency information may be utilized, as will be explained with reference to FIG. 5.

(47) In Ag. 3, a scenario is illustrated in which the first physical resource **210** is selected for execution of the first part **308** and the second physical resource **212** is selected for execution of the second part **310**. In response to such a selection, in an example, the first scheduler **220** may map the first logical identifier **328** to the identifier of the first physical resource **210** and map the second logical identifier **330** to the identifier of the second physical resource **212**. If each logical identifier is to correspond to a computational unit, the first scheduler **220** may map the first logical identifier **328** to the identifier of a computational unit of the first physical resource **210** and map the second logical identifier **330** to the identifier of a computational unit of the second physical resource **212**. Further, the mappings may be provided to the first loader **332**.

(48) Upon receiving the mappings, the first loader **332** may replace the first logical identifier **328** in the neural network **306** with the identifier of the first physical resource **210** (or with the identifier of the computational unit of the first physical resource **210**) and replace the second logical identifier **330** in the neural network **306** with the identifier of the second physical resource **212** (or with the identifier of the computational unit of the second physical resource **212**). Thus, all references to logical identifiers in the neural network **306** may be replaced with the identifiers of the physical resources (or computational units) with which the identifiers were mapped by the first scheduler **220**. As mentioned earlier, an identifier of a physical resource/computational unit may be a physical address of the physical resource/computational unit. In addition, the first loader **332** may load a compute block (which includes a part of the neural network **306** and its associated neural network parameters) onto a physical resource to which the part is assigned. The first loader **332** may perform such a loading function onto those physical resources that are present in the first computing device **204**. For example, of the physical resources selected, the first physical resource **210** is present in the first computing device **204**. Accordingly, the first loader **332** may load the first compute block **324** onto the first physical resource **210**. To load the second compute block **326** onto the second physical resource **212**, in an example, the first scheduler **220** may provide the second compute block **326** to a second scheduler **334**, which is present in the second computing device **206**. The second scheduler **334** may then load the second compute block **326** onto the second physical resource **212** through a second loader (not shown in FIG. 3) of the second computing device **206**.

(49) In an example, the first scheduler **220** may designate one of the selected physical resources as a conductor resource and the other selected physical resources as member resources. The conductor resource may be responsible for execution of one part of the neural network **306**. In addition, the conductor resource may be responsible for receiving the final output of execution of the neural network **306**, such as from the physical resource that executes the final part of the neural network, and for outputting the final output. In an example, the first scheduler **220** may designate a physical resource that is part of the first computing device **204** (which received the instruction to execute the neural network **306**) as the conductor resource. For example, if the first physical resource **210** and the second physical resource **212** are selected for execution of the neural network **306**, the first scheduler **220** may designate the first physical resource **210** as the conductor resource. Such a designation may allow the first computing device **204**, which received the instruction to execute the neural network **306**, to provide the output of execution of the neural network **306**. In an example, if the first computing device **204** does not have an available physical resource for execution of a part (e.g., if the first physical resource **210** is busy executing a part of another neural network), the first scheduler **220** may output an error indicating that the execution of the neural network **306** cannot be initiated and managed by the first computing device **204**. In such a case, the user may provide the instruction to execute the neural network **306** to another computing device, such as the second computing device **206**, or may provide the instruction to the first computing device **204** at a later point of time.

(50) Although the first computing device **204** and the second computing device **206** are each shown as having one physical resource, the computing devices may have more than one physical resource. Further, the physical resources may be connected among themselves through various communication links.

(51) FIG. 4 illustrates the connections among the physical resources of the first computing device **204** and the second computing device **206**, according to an example implementation of the present subject matter. As illustrated, in addition to the first physical resource **210**, the first computing device **204** may include a sixth physical resource **402** and a seventh physical resource **404**. Further, the second computing device **206** may include the third physical resource **214**, in addition to the second physical resource **212**.

(52) The physical resources in a computing device may be connected with each other through any of a variety of communication links. For example, a first communication link **406** connecting the first physical resource **210** and the sixth physical resource **402**, a second communication link **408** connecting the sixth physical resource **402** and the seventh physical resource **404**, and a third communication link **410** connecting the second physical resource **212** and the third physical resource **214** may be of a first type. Further, a fourth communication link **412** connecting the first physical resource **210** and the seventh physical resource **404** may be of a second type. Further, physical resources of different computing devices may also be connected through a communication link. For instance, a fifth communication link **414** may connect the first physical resource **210** and the second physical resource **212**. The fifth communication link **414** may be of a third type. The fifth communication link **414** may be a dedicated connection between the first physical resource **210** and the second physical resource **212** and may be different than a communication link, such as an ethernet connection, between the first computing device **204** and the second computing device **206**. In an example, the communication between physical resources on different computing devices may not happen over the communication link connecting the computing devices, and may happen over the dedicated communication link connecting the physical resources.

(53) A speed of one type of communication link may be higher than that of another type of communication link. For example, the first type of communication link may have a higher speed (and therefore, a lower latency) as compared to the second type of communication link and the second type of communication link may have a higher speed as compared to the third type of communication link. The various types of communication link may include an Infiniband link and

a PCI Express (PCIe) link. In an example, two physical resources may be connected by more than one communication link, each of a different type. For instance, in addition to the first communication link **406**, such as an Infiniband link, another communication link (not shown in FIG. **4**) that is of the second type, such as a PCIe link, may connect the first physical resource **210** and the sixth physical resource **402**.

(54) In an example, any two physical resources within a computing device may be directly connected to each other, i.e., such that there is no third physical resource in a connection between the two physical resources. For instance, as illustrated, the seventh physical resource **404** is directly connected to both the remaining physical resources in the first computing device **204**. However, some physical resources in different computing devices may not be connected to each other directly. For instance, as illustrated, the third physical resource **214** is indirectly connected to the first physical resource **210**, through the second physical resource **212**. Similarly, the second physical resource **212** is connected to the sixth physical resource **402** and the seventh physical resource **404** indirectly, through the first physical resource **210**. The communication between two indirectly-connected physical resources may happen through an intermediate physical resource that connects the indirectly-connected physical resources. The indirect connection between two physical resources may increase the latency in communication between the physical resources due to the involvement of the intermediate physical resource in the communication.

(55) The latency in communication between two physical resources may be considered by the first scheduler **220** for scheduling parts of the neural network **306** to the physical resources. For instance, the first scheduler **220** may not assign the first part **308** and the second part **310** for execution to the seventh physical resource **404** and the third physical resource **214** respectively, as the high latency that would be involved in communication of data between the seventh physical resource **404** and the third physical resource **214** may reduce the speed of execution of the neural network **306**. To schedule the parts of the neural network based on the latency, the first scheduler **220** may utilize duster latency information, as will be explained below.

(56) FIG. **5** illustrates a cluster latency table **500** providing information of latency between physical resources of computing devices, according to an example implementation of the present subject matter. The cluster latency table **500** may also be referred to as the latency table **500**. The information provided by the latency table **500** may be referred to as cluster latency information. For the sake of simplicity, the latency table **500** is shown in a case where the cluster **202** has two computing devices (the first and the second computing devices **204** and **206**) and where the physical resources of the first computing device **204** and the physical resources of the second computing device **206** are connected in the manner as illustrated in FIG. **4**. However, it will be understood that the latency table **500** may be adapted for any number of computing devices and any manner of connections between the physical resources.

(57) The latency table **500** may provide information of latency in communication between each physical resource of the cluster **202** and each remaining physical resource of the cluster **202**. For example, the latency table **500** indicates the latency between the first physical resource **210** and each remaining physical resource in the cluster **202**, i.e., the second physical resource **212**, the third physical resource **214**, the sixth physical resource **402**, and the seventh physical resource **404**. The latency in communication between two physical resources may depend on the type of communication link between them. Accordingly, the latency values in the latency table **500** may be higher for a slower communication link. For instance, the latency value for the first communication link **406** (between the first physical resource **210** and the sixth physical resource **402**) ('10') may be less than that for the fifth communication link **414** (between the first physical resource **210** and the second physical resource **212**) ('30'). In an example, a ratio of two latency values may provide a measure of the slowness of one communication link as compared to another. For instance, a latency value of 10 for the first communication link **406** and a latency value of 30 for the fifth communication link **414** may indicate that the first communication link **406** is thrice as fast as the

fifth communication link **414**. In some cases, two physical resources may be connected by more than one communication link, each of a different type. In such cases, the latency value indicated in the table may be that corresponding to the higher-speed communication link. For example, even if the first physical resource **210** and the sixth physical resource **402** are connected to each other by communication links of the first type and the second type, the latency value indicated in the latency table **500** may be that of the first type.

(58) In an example, the latency value assigned to an indirect connection, such as the connection between the first physical resource **210** and the third physical resource **214**, may be an arbitrary value. Such a value may be higher than the latency values assigned for direct communication links (e.g., the fifth communication link **414** and the third communication link **410**), to indicate that the communication between the two concerned devices is likely to be slower than that between two directly-connected devices. In an example, the arbitrary value may be 50.

(59) The latency table **500** may be generated, for example, by the first computing device **204**, such as by the first scheduler **220** of the first computing device **204**. To generate the latency table **500**, the first computing device **204** may determine the communication links between the physical resources of the first computing device **204** and accordingly populate the latency values. In an example, the determination of the links between the physical resources may be performed by the first loader **332**. Further, the first loader **332** may provide the result of the determination to the first scheduler **220**. Further, the first computing device **204** may query the second computing device **206** for the various physical resources in the second computing device **206** and the latency in communication between the physical resources or the communication links between those physical resources, and accordingly populate latency values. The information of the physical resources in the second computing device **206** and the communication links between those physical resources may be provided, for example, by the second scheduler **334**. Further, the first computing device **204** may determine the communication links existing between physical resources of the first computing device **204** and physical resources of the second computing device **206** and accordingly populate latency values in the latency table **500**.

(60) In an example, when the third computing device **208** (not shown in FIG. 5), is to join the cluster **202**, the third computing device **208** may send a request for joining to the first computing device **204**. Upon receiving the request, the first computing device **204** may receive information of the physical resources in the third computing device **208** and the connections between those physical resources and update the latency table **500** accordingly. The information of the connections may be, for example, the information of communication links between the physical resources or the information of latency in communication between the physical resources. The first computing device **204** may also determine the connections between the physical resources of the first computing device **204** and that of the third computing device **208** and determine the connections between the physical resources of the second computing device **206** and that of the third computing device **208** and accordingly update the latency table **500**. The updated latency table may also be shared with the second computing device **206** and the third computing device **208**, so that all computing devices of the cluster **202** may have the latest information regarding the computing devices in the cluster **202**, their respective physical resources, and the latencies that would be involved in communication between any two physical resources of the cluster **202**.

(61) In another example, upon receiving a request to join from the third computing device **208**, the second computing device **206** may update the latency table **500** independently of the first computing device **204**. The information of the physical resources in the cluster **202** and the connections between the physical resources may be referred to as cluster topology information.

(62) The first scheduler **220** may utilize the latency table **500** to schedule the parts of the neural network **306**. To perform the scheduling, the first scheduler **220** may first determine the physical resources of the cluster **202** that are available for execution of neural network parts, such as the physical resources that are not occupied with execution of parts of other neural networks. The first

scheduler **220** may also receive the information of the number of parts in the neural network **306**.

(63) In an example, to perform the scheduling, the first scheduler **220** may initially perform a tentative scheduling, according to which the first scheduler **220** tentatively assigns the parts to any of the available physical resources and may compute the total latency that would be involved in communication of data among the assigned physical resources for execution of the neural network **306**. For example, consider that all physical resources of the cluster **202** are available for execution of parts of the neural network **306** and that the neural network **306** includes the first part **308**, the second part **310**, and a third part (not shown in FIG. 5). Consider also that output from execution of the first part **308** is to be provided to the second part **310** and that output from execution of the second part **310** is to be provided to the third part. Further, consider that output from execution of the third part is the final output and is therefore to be provided to the physical resource that would act as the conductor resource, for being shared as the output of execution of the neural network **306**. In such a case, the first scheduler **220** may tentatively assign the executions of the first part **308** to the first physical resource **210**, the second part **310** to the second physical resource **212**, and the third part to the third physical resource **214**. The first scheduler **220** may also tentatively designate the first physical resource **210** as the conductor resource. Accordingly, the total latency value may be computed as a sum of (i) a latency value involved in communicating between the first physical resource **210** and the second physical resource **212** (for communicating output from execution of the first part **308** to the second part **310**), (ii) a latency value involved in communicating between the second physical resource **212** and the third physical resource **214** (for communicating output from execution of the second part **310** to the third part), and (iii) a latency value involved in communicating between the third physical resource **214** and the first physical resource **210** (for communicating output from execution of the third part to the conductor resource). Thus, the total latency value for such an assignment is 90 (30+10+50). In this manner, the first scheduler **220** may determine the latency values for various tentative assignments of physical resources and various tentative designations of conductor resources.

(64) Subsequently, the first scheduler **220** may select the tentative assignment and tentative designation that involves the least latency value for scheduling the parts of the neural network **306**. For instance, first scheduler **220** may assign the first part **308** to the first physical resource **210**, the second part **310** to the sixth physical resource **402**, and the third part to the seventh physical resource **404** and may designate the first physical resource **210** as the conductor resource, since such an assignment provides the least total latency value (10+10+20=40).

(65) In an example, the tentative assignment for one part may depend on the physical resource to which the previous part is tentatively assigned. For instance, the physical resource to which a particular part is tentatively assigned may be the one having the least latency for communication with the physical resource to which the previous part is tentatively assigned. This will be explained with the help of an example. Consider that the first part **308** is tentatively assigned to the first physical resource **210**. Accordingly, the second part **310** may be tentatively assigned to the sixth physical resource **402**, as among the remaining available physical resources, the sixth physical resource **402** has the least latency for communication with the first physical resource **210**. Similarly, the third part may be tentatively assigned to the seventh physical resource **404**, as among the remaining available physical resources, the seventh physical resource **404** has the least latency in communication with the sixth physical resource **402**.

(66) In both the above examples, the first scheduler **220** may assign the parts to the physical resources such that the latency that would be caused due to communication of data among parts of the neural network **306** is as small as possible. Thus, the present subject matter achieves a quick execution of the neural network **306**.

(67) In an example, the parts of the neural network **306** may be assigned to computational units in the physical resources. Further, the parts may have to be assigned to physical resources that are indirectly connected to each other due to unavailability of directly-connected physical resources.

The scheduling in such scenarios will be explained with reference to FIG. 6. Although in FIG. 5, the cluster latency information is expressed in the form of a table, in other examples the information may be expressed in other forms, such as in the form of a weighted graph.

(68) FIG. 6 illustrates the cluster **202** having the first computing device **204** and the second computing device **206**, according to an example implementation of the present subject matter. In an example, the physical resources in the computing devices may be DPEs. For example, the first computing device **204** includes a first DPE **602**, a second DPE **604**, and a third DPE **606**, which correspond to the first physical resource **210**, the sixth physical resource **402**, and the seventh physical resource **404** respectively. Similarly, the second computing device **206** may include a fourth DPE **608** and a fifth DPE **610**, which correspond to the second physical resource **212** and the third physical resource **214** respectively. A DPE may perform a MVM operation as a dot-product computation. A DPE may include one or more computation units, which are also referred to as tiles. For example, the first DPE **602** may include a first tile **612**, the fourth DPE **608** may include a second tile **614**, and the fifth DPE **610** may include a third tile **616** and a fourth tile **618**. Similarly, the second DPE **604** and the third DPE **606** may each include one or more tiles (not shown in FIG. 6).

(69) In an example, each tile may include a memory and a plurality of cores (not shown in FIG. 6). The memory may be loaded with the instructions that are to be executed by the tile. A core may include a plurality of MVM units, which can perform MVM operations involved in execution of a neural network. In an example, a MVM unit may be programmed with a matrix, such as a weight matrix, that is involved in a MVM operation. Further, the MVM unit may receive the vector that is to be multiplied with the matrix as an input and perform the MVM operation. In an example, in addition to the MVM units, each core may include units that can perform other operations, such as vector-vector addition operation, scalar-scalar addition operation, and matrix-matrix addition, that are involved in execution of a neural network.

(70) The DPEs of the cluster **202** may be connected in various ways. For example, as illustrated, the first DPE **602**, the second DPE **604**, and the third DPE **606** may be directly connected between each other. Similarly, the fourth DPE **608** and the fifth DPE **610** may be directly connected to each other. However, the fifth DPE **610** may be indirectly connected to the first DPE **602** through the fourth DPE **608**.

(71) The neural network **306** may include the first part **308**, the second part **310**, and a third part **620**. In an example, an output of execution of the first part **308** is to be provided as an input to the second part **310**. Similarly, an output of execution of the second part **310** may have to be provided as an input to the third part **620**. The third part **620** may be the final part of the neural network **306**. Accordingly, the output from execution of the third part **620** may have to be provided to the DPE that executes the first part **308**, for providing the output of execution of the neural network **306**. A part that is to receive an output of execution of another part and a part that is to provide its output to another part may be referred to as being related to the other part. Accordingly, the first part **308**, the second part **310**, and the third part **620** are related to each other.

(72) The first scheduler **220** may assign the various parts of the neural network **306** to the tiles of the DPEs in the cluster **202** based on availability of tiles and based on latency in communication between the DPEs. In an example, the first scheduler **220** may attempt to schedule the parts such that a part that is to receive the output of execution of another part is assigned to the same DPE as the other part or to a DPE that is directly connected to the DPE to which the other part is assigned. For example, the first scheduler **220** may attempt to assign both the first part **308** and the second part **310** to the same DPE (i.e., to different tiles of the same DPE). If such an assignment is not possible, the first scheduler **220** may attempt to assign the first part **308** and the second part **310** to DPEs that are directly connected to each other, such as the first DPE **602** and the second DPE **604**. The first scheduler **220** may attempt to perform such an assignment because assignment of related parts, such as the first part **308** and the second part **310**, to indirectly-connected DPEs increases

latency in communication of data from one part to another, as explained above. Further, the first scheduler **220** may attempt to schedule such that the DPE to which the final part of a neural network is assigned is same as or is directly connected to the DPE to which the first part of the neural network is assigned. For instance, the first scheduler **220** may attempt to assign the first part **308** and the third part **620** to the same DPE. If such an assignment is not possible, the first scheduler **220** may attempt to assign the first part **308** and the third part **620** to directly-connected DPEs. Such an assignment may be attempted because, the DPE to which the first part of the neural network is assigned is to be the conductor resource and is to receive output from execution of the final part, so that the conductor resource can generate the output of execution of the neural network. However, in some cases, the above-explained attempts may fail. One such case is explained below.

(73) Consider that tiles of the second DPE **604** and of the third DPE **606** are unavailable for execution of parts of the neural network **306**, such as due to execution of parts of another neural network. Consider also that the first DPE **602** has a single tile, i.e., the first tile **612**, or that other tiles (not shown in FIG. **6**) of the first DPE **602** are unavailable. In such a case, the tiles of the cluster **202** that are available for execution may be the first tile **612** and second through fourth tiles **614-618**. Accordingly, the three parts of the neural network **306** are to be assigned among the four tiles. Further, the first part **308** is to be assigned to the first tile **612**, as the first DPE **602** is to be the conductor resource (because the conductor resource is to be part of the computing device that received the instruction to execute the neural network **306**, i.e., the first computing device **204**, and the first DPE **602** is the only available DPE in the first computing device **204**) and the first tile **612** is the only (or only available) tile in the first DPE **602**. Thus, the remaining two parts are to be assigned among the second through fourth tiles **614-618**. Thus, at least one part of the neural network **306** is to be assigned to the fifth DPE **610**, and such an assignment involves a communication between the first DPE **602** and the fifth DPE **610**.

(74) In cases where it is not possible to assign related parts to directly-connected DPEs, such as in the one explained above, the first scheduler **220** may assign the parts to indirectly-connected DPEs. Even while performing such an assignment, the first scheduler **220** may attempt to keep the latency as less as possible. For instance, the first scheduler **220** may attempt to assign the third part **620** to a DPE that is connected to a DPE to which the first part **308** is assigned through a single intermediate DPE. Further, while performing such an assignment, the first scheduler **220** may reserve a portion, such as one tile, of the intermediate DPE that connects the two DPEs for transferring data between the two DPEs. For instance, referring to the above example, the first scheduler **220** may assign execution of the first part **308** to the first tile **612**, execution of the second part **310** to the third tile **616**, and execution of the third part **620** to the fourth tile **618**. Further, the second tile **614**, which is part of the fourth DPE **608** that connects the first DPE **602** and the fifth DPE **610**, may be reserved for transferring output of execution of the third part **620** to the first DPE **602**.

(75) As explained above, the first scheduler **220** attempts to assign related parts to the same/neighboring DPEs and to assign the first and final parts to the same/neighboring DPEs and assigns accordingly if the attempts succeed. Thus, the scheduling according to the present subject matter achieves least possible latency in execution of a neural network in a clustered environment. Even if such a scheduling is not possible, by scheduling the parts to indirectly-connected DPEs and by reserving a portion of an intermediate DPE for communication of data between the indirectly-connected DPEs, the present subject matter ensures that execution of a neural network is achieved regardless of the cluster topology and regardless of DPE availability.

(76) FIGS. **7** and **8** illustrate methods **700** and **800**, respectively, for causing execution of neural networks, according to example implementations of the present subject matter. The orders in which the methods **700** and **800** are described is not intended to be construed as a limitation, and any number of the described method blocks may be combined in any order to implement the methods **700** and **800**, or alternative methods. Furthermore, the methods **700** and **800** may be implemented

by processing resource(s) or computing device(s) through any suitable hardware, non-transitory machine-readable instructions, or a combination thereof.

(77) It may be understood that blocks of the methods **700** and **800** may be performed by programmed computing devices and may be executed based on instructions stored in a non-transitory computer readable medium. The non-transitory computer readable medium may include, for example, digital memories, magnetic storage media, such as magnetic disks and magnetic tapes, hard drives, or optically readable digital data storage media. In an example, the methods **700** and **800** may be performed by a computing device, such as the computing device **100** or the first computing device **204**. Although the methods **700** and **800** may be implemented in a variety of systems, the methods **700** and **800** are described in relation to the first computing device **204**, for ease of explanation.

(78) Referring to FIG. 7, at block **702**, a neural network including a first number of parts may be received by a first computing device of a duster. Each part is to be executed by a physical resource. The neural network may be, for example, the neural network **306**, the first computing device may be the first computing device **204**, and the duster may be the duster **202**. Further, the physical resource may be, for example, a DPE.

(79) At block **704**, cluster latency information may be received. The cluster latency information may indicate a plurality of physical resources present in computing devices of the cluster and latency in communication among the plurality of physical resources. The cluster latency information may be, for example, the information included in a latency table, such as the latency table **500**. Accordingly, the receipt of the cluster latency information may involve receiving the latency table.

(80) At block **706**, a first number of physical resources that are to execute the first number of parts may be determined. The first number of physical resources may be part of the plurality of physical resources. The determination of the first number of physical resources may be based on a latency in communication among the physical resources, as explained above. To determine the first number of physical resources based on the latency, the cluster latency information may be utilized.

(81) In an example, the neural network may include a first part and a second part, such as the first part **308** and the second part **310** respectively. Further, an output from execution of the first part may have to be sent to the second part. In such a case, determining the first number of physical resources may include determining a physical resource that is to be selected for execution of the second part if a first physical resource is selected for execution of the first part, based on a latency in communication between the first physical resource and other physical resources of the plurality of physical resources. For instance, as explained above, a first physical resource may be tentatively selected for execution of the first part. In response to the tentative selection, a physical resource to which the second part is to be tentatively assigned may be determined based on latencies involved in communication of the other physical resources with the first physical resource.

(82) In an example, as explained with reference to FIG. 6, the first computing device may attempt to assign both the first part and the second part to the same physical resource, as communication of output of the first part to the second part may be performed without a significant latency in that case. If such an attempt fails, the first computing device may attempt to assign the second part to a physical resource that is directly connected to the physical resource to which the first part is assigned. As explained above, a direct connection between two physical resources may refer to a connection that does not involve an intermediate physical resource, to facilitate communication of data between the two physical resources.

(83) Further, in an example, an output from execution of a third part of the neural network is to be sent to a physical resource that is to execute a first part. The output from the third part may have to be utilized by the physical resource that is to execute the first part to generate an output of execution of the neural network. The first part may be, for example, the part that is to be executed at the beginning, such as before execution of the other parts. The third part may be, for example, a

final part that is to be executed after execution of the remaining parts. The first part may be, for example, the first part **308** and the third part may be, for example, the third part **620**. The physical resource that is to execute the third part may have to communicate the output of the third part to the physical resource that executes the first part. Accordingly, the method **700** may include determining a physical resource that is to be selected for execution of the third part if a first physical resource is selected for execution of the first part, based on a latency in communication between the first physical resource and a plurality of physical resources. For example, upon tentatively assigning execution of the first part to the first physical resource, the first computing device may determine if execution of the third part may also be assigned to the first physical resource. If that is not possible, such as due to tentative assignment of intermediate parts to the first physical resource, the first computing device may determine if execution of the third part may be assigned to a physical resource that is directly connected to the first physical resource.

(84) At block **708**, the first computing device may cause execution of the parts of the neural network by the first number of physical resources. To cause the execution, in an example, the first computing device may load the parts that are assigned to physical resources in the first computing device onto those physical resources. For instance, if the first part of the neural network is assigned to a first physical resource, such as the first physical resource **210**, in the first computing device, the first computing device may load the first part onto the first physical resource. Further, the first computing device may transmit parts that are assigned to physical resources in other computing devices to those computing devices. For instance, if the second part of the neural network is assigned to a second physical resource, such as the second physical resource **212**, in the second computing device, the first computing device may transmit the second part to the second physical resource. Subsequently, the second computing device may load the second part onto the second physical resource.

(85) FIG. **8** illustrates a method **800** for scheduling parts of a neural network to physical resources, according to an example implementation of the present subject matter. As explained above, output from the third part of the neural network may have to be sent to the physical resource executing the first part to allow generation of an output of execution of the neural network. At block **802**, it is determined whether execution of the third part and execution of the first part can be assigned to physical resources that are directly connected to each other. If yes, at block **804**, the third part and the first part are assigned to physical resources that are directly connected to each other.

(86) However, in some cases, such an assignment may not be possible. This may be, for example, because the number of parts in the neural network may be more than the number of available computation units in directly-connected physical resources, as exemplified in FIG. **6** and its corresponding description. In response to a determination that the execution of the third part and execution of the first part are not assignable to physical resources that are directly connected to each other, at block **806**, the first part may be assigned to a first physical resource and, at block **808**, the third part may be assigned to a third physical resource that is indirectly connected to the first physical resource. The third physical resource may be connected to the first physical resource indirectly through a second physical resource. The first physical resource, the second physical resource, and the third physical resource may be, for example, the first DPE **602**, the fourth DPE **608**, and the fifth DPE **610** respectively.

(87) Further, at block **810**, a portion of the second physical resource may be reserved for transferring the output of execution of the third part to the first physical resource. The portion may include, for example, a tile, such as the second tile **614**.

(88) FIG. **9** illustrates a computing environment **900** implementing a non-transitory computer-readable medium for causing execution of neural networks, according to an example implementation of the present subject matter. In an example, the non-transitory computer-readable medium **902** may be utilized by a computing device, such as the computing device **100** or the first computing device **204**. In an example, the computing environment **900** may include a processing

resource **904** communicatively coupled to the non-transitory computer-readable medium **902** through a communication link **906**. The processing resource **904** may be, for example, the processor **102**.

(89) The non-transitory computer-readable medium **902** may be, for example, an internal memory device or an external memory device. In an example, the communication link **906** may be a direct communication link, such as any memory read/write interface. In another example, the communication link **906** may be an indirect communication link, such as a network interface. In such a case, the processing resource **904** may access the non-transitory computer-readable medium **902** through a network **908**. The network **908** may be a single network or a combination of multiple networks and may use a variety of different communication protocols.

(90) The processing resource **904** and the non-transitory computer-readable medium **902** may also be communicatively coupled to computing device(s) **910**, which may be part of a cluster of computing devices, such as the cluster **202**.

(91) In an example implementation, the non-transitory computer-readable medium **902** includes a set of computer-readable instructions for causing execution of a neural network. The set of computer-readable instructions can be accessed by the processing resource **904** through the communication link **906** and subsequently executed.

(92) Referring to FIG. 9, in an example, the non-transitory computer-readable medium **902** includes instructions **914** that cause the processing resource **904** to receive a neural network including a first logical identifier and a second logical identifier. The first logical identifier corresponds to a first part of the neural network and the second logical identifier corresponds to a second part of the neural network. The neural network, the first part, the second part, the first logical identifier, and the second logical identifier may be, for example, the neural network **306**, the first part **308**, the second part **310**, the first logical identifier **328**, and the second logical identifier **330** respectively.

(93) In an example, each of the first part and the second part includes a set of neural network instructions, also referred to as a set of instructions. The set of instructions may include instructions to perform matrix-vector multiplications, instructions to perform vector-vector additions, instructions to perform loading of data, instructions to perform storing of data, instructions to perform sending of data, or a combination thereof. Each of the first part and the second part may be associated with weights, biases, or both.

(94) The non-transitory computer-readable medium **902** may include instructions **916** that cause the processing resource **904** to determine that a first physical resource in a first computing device is to execute the first part and that a second physical resource in a second computing device is to execute the second part. Such a determination may be based on a latency in a communication link between the first physical resource and the second physical resource, as explained above. The first computing device and the second computing device may be part of a cluster, such as the cluster **202**.

(95) In an example, to determine the latency in communication between first physical resource and the second physical resource, cluster latency information may be utilized. The cluster latency information may indicate latency in communication between each physical resource in the cluster and remaining physical resources in the cluster. The cluster latency information may be, for example, the information included in the latency table **500**. In an example, the cluster latency information may be generated by querying the second computing device for latency in communication between physical resources of the second computing device and by determining latency in communication between physical resources of the first computing device and physical resources of the second computing device.

(96) In response to the determination by the instructions **916**, instructions **918** of the non-transitory computer-readable medium **902** cause the processing resource **904** to replace the first logical identifier in the neural network with an identifier of the first physical resource to cause execution of

the first part by the first physical resource. Further, instructions **920** of the non-transitory computer-readable medium **902** cause the processing resource **904** to replace the second logical identifier in the neural network with an identifier of the second physical resource to cause execution of the second part by the second physical resource.

(97) In an example, each of the first part and the second part is executable by a tile of a physical resource. Further, the identifier of the first physical resource may be a first physical address of a first tile of the first physical resource, and the identifier of the second physical resource may be a second physical address of a second tile of the second physical resource. The first tile may be, for example, the first tile **612** and the second tile may be, for example, the third tile **616**.

(98) The present subject matter allows execution of a neural network in a cluster of computing devices, where each computing device has one or more physical resources for execution of neural network parts. Thus, the number of physical resources available for execution of a neural network increases. Accordingly, the present subject matter may prevent reprogramming or reduce the number of reprogramming of neural network data on physical resources (e.g., programming weight matrices on MVM units of DPEs), which may have to be performed if the number of physical resources available is less than the number of neural network parts. For example, using techniques of the present subject matter, weight matrices involved in execution of the neural network may remain programmed on the physical resources during the entire period of execution.

(99) The present subject matter can be utilized for neural network execution in a computing environment in which the computing devices have a small number of physical resources. Thus, data is not to be transmitted from the computing environment to a cloud environment for execution of the neural networks, thereby eliminating consumption of network resources latency that would otherwise be involved in transmission of data to and from the cloud environment.

(100) The present subject matter provides an efficient manner of scheduling neural network parts on the physical resources by considering data flow involved in the execution of the neural network and the latency involved communication of data among physical resources. Thus, execution of the neural network may be made faster.

(101) The techniques of the present subject matter can be implemented in various types of cluster topologies, where physical resources are connected among them in a variety of configurations. Further, the manner in which the neural network model is compiled by the compiler may not have to be changed depending on the cluster topology. Thus, the compiler may be kept agnostic to the underlying cluster topology. Further, an application developer developing the neural network model is not to be aware of the cluster topology.

(102) Although implementations of execution of neural networks have been described in language specific to structural features and/or methods, it is to be understood that the present subject matter is not necessarily limited to the specific features or methods described. Rather, the specific features and methods are disclosed and explained as example implementations.

Claims

1. A computing device comprising: a processor; and a memory comprising instructions executable by the processor to: receive a neural network to be executed by a plurality of hardware accelerators, the neural network comprises a first logical identifier corresponding to a first part and a second logical identifier corresponding to a second part, the first and second logical identifiers acting as placeholders for identifiers of hardware accelerators executing, respectively, the first and second parts of the neural network; determine data-transfer latencies among the plurality of hardware accelerators based on speeds of communication links among the hardware accelerators; identify first and second hardware accelerators having least data-transfer latency among the plurality of hardware accelerators, wherein the identified first and second hardware accelerators belong, respectively, to first and second computing devices, which are part of a cluster; replace the first

logical identifier in the neural network with an identifier of the identified first hardware accelerator; replace the second logical identifier in the neural network with an identifier of the identified second hardware accelerator; and execute the neural network, the first part being executed by the identified first hardware accelerator and the second part being executed by the identified second hardware accelerator.

2. The computing device of claim 1, wherein, to determine data-transfer latency between the first hardware accelerator and the second hardware accelerator, the instructions are executable by the processor to utilize cluster latency information that indicates data-transfer latency between each hardware accelerator in the cluster and remaining hardware accelerators in the cluster.

3. The computing device of claim 2, wherein the computing device is same as the first computing device, wherein the instructions are executable by the processor to generate the cluster latency information, and wherein, to generate the cluster latency information, the instructions are executable to: query the second computing device for data-transfer latency between hardware accelerators within the second computing device; and determine data-transfer latency between hardware accelerators within the first computing device and hardware accelerators within the second computing device.

4. The computing device of claim 2, wherein the instructions are executable by the processor to determine data-transfer latency between the first hardware accelerator and another hardware accelerator in the first computing device.

5. The computing device of claim 2, wherein the instructions are executable by the processor to: receive a request from a third computing device to join the cluster; and update the cluster latency information with information of: hardware accelerators in the third computing device; data-transfer latency between the hardware accelerators in the third computing device; and data-transfer latency between the first hardware accelerator and each hardware accelerator in the third computing device.

6. The computing device of claim 1, wherein each hardware accelerator is a dot-product engine (DPE) having a plurality of tiles.

7. The computing device of claim 6, comprising instructions executable to generate the neural network by compiling a neural network model, wherein, to generate the neural network, the instructions are executable to generate the first part and the second part such that each of the first part and the second part is executable by a tile of a DPE.

8. A method comprising: receiving, by a first computing device of a cluster, a neural network comprising a first number of parts, wherein each part is to be executed by a hardware accelerator, and wherein each part includes a logical identifier acting as a placeholder for an identifier of a hardware accelerator executing the part; receiving, by the first computing device, cluster latency information indicating a plurality of hardware accelerators present in computing devices of the cluster and data-transfer latencies among the plurality of hardware accelerators, the data-transfer latencies being determined based on speeds of communication links among the hardware accelerators; identifying, by the first computing device, a first number of hardware accelerators with least data-transfer latencies among them from the plurality of hardware accelerators; replacing logical identifiers in the first number of parts with identifiers associated with the identified first number of hardware accelerators; and executing the first number of parts of the neural network by the first number of hardware accelerators.

9. The method of claim 8, wherein an output from execution of a first part of the neural network is to be sent to a second part of the neural network and wherein determining the first number of hardware accelerators comprises: determining a hardware accelerator that is to be selected for execution of the second part if a first hardware accelerator is selected for execution of the first part, based on data-transfer latency between the first hardware accelerator and other hardware accelerators of the plurality of hardware accelerators.

10. The method of claim 8, wherein an output from execution of a third part of the neural network is to be sent to a hardware accelerator that is to execute a first part of the neural network, to cause

the hardware accelerator that is to execute the first part to generate an output of execution of the neural network, and wherein determining the first number of hardware accelerators comprises: determining a hardware accelerator that is to be selected for execution of the third part if a first hardware accelerator is selected for execution of the first part based on data-transfer latency in between the first hardware accelerator and a plurality of hardware accelerators.

11. The method of claim 10, wherein determining the first number of hardware accelerators comprises: determining whether execution of the third part and execution of the first part are assignable to hardware accelerators that are directly connected to each other; in response to a determination that the execution of the third part and execution of the first part are not assignable to hardware accelerators that are directly connected to each other, determining the first number of hardware accelerators comprises: determining that execution of the first part is to be assigned to a first hardware accelerator; determining that execution of the third part is to be assigned to a third hardware accelerator, wherein the third hardware accelerator is connected to the first hardware accelerator indirectly through a second hardware accelerator; and reserving a portion of the second hardware accelerator for transferring the output of execution of the third part to the first hardware accelerator.

12. The method of claim 11, wherein the second hardware accelerator is a dot-product engine (DPE) and wherein the portion of the second hardware accelerator reserved comprises a tile of the DPE.

13. The method of claim 8, wherein the first number of hardware accelerators comprises a first hardware accelerator in the first computing device and wherein executing the first number of parts comprises: loading, by the first computing device, a first part of the neural network onto the first hardware accelerator; and transmitting, by the first computing device, a second part of the neural network to a second computing device for loading of the second part onto a second hardware accelerator of the second computing device.

14. The method of claim 13, comprising loading, by the second computing device, the second part of the neural network onto the second hardware accelerator.

15. A non-transitory computer-readable medium comprising instructions, the instructions being executable by a processing resource to: receive a neural network comprising a first logical identifier and a second logical identifier, wherein the first logical identifier corresponds to a first part of the neural network and the second logical identifier corresponds to a second part of the neural network, and wherein the first and second logical identifiers act as placeholders for identifiers of hardware accelerators executing, respectively, the first and second parts of the neural network; determine data-transfer latencies among a plurality of hardware accelerators based on speeds of communication links among the hardware accelerators; identify first and second hardware accelerators having least data-transfer latency among the plurality of hardware accelerators, wherein the identified first and second hardware accelerators belong, respectively, to the first and second computing devices, which are part of a cluster; replace the first logical identifier in the neural network with an identifier of the identified first hardware accelerator; replace the second logical identifier in the neural network with an identifier of the identified second hardware accelerator; and execute the neural network, the first part being executed by the identified first hardware accelerator and the second part being executed by the identified second hardware accelerator.

16. The non-transitory computer-readable medium of claim 15, wherein, to determine data-transfer latency between first hardware accelerator and the second hardware accelerator, the instructions executable by the processing resource to utilize cluster latency information that indicates data-transfer latency between each hardware accelerator in the cluster and remaining hardware accelerators in the cluster.

17. The non-transitory computer-readable medium of claim 16, wherein the instructions are executable by the processing resource to generate the cluster latency information by: querying the

second computing device for data-transfer latency between hardware accelerators of the second computing device; and determining data-transfer latency between hardware accelerators of the first computing device and hardware accelerators of the second computing device.

18. The non-transitory computer-readable medium of claim 15, wherein each of the first part and the second part comprises a set of neural network instructions, the set of neural network instructions comprising instructions to perform one or more of: matrix-vector multiplications, vector-vector additions, loading of data, storing of data, and sending of data, and wherein each of the first part and the second part is associated with one or more of: weights and biases.

19. The non-transitory computer-readable medium of claim 15, wherein each hardware accelerator is a dot-product engine (DPE) having a plurality of tiles, wherein each of the first part and the second part is executable by a tile of a DPE, wherein the identifier of the first hardware accelerator is a first physical address of a first tile, and wherein the identifier of the second hardware accelerator is a second physical address of a second tile.
