

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12406209
Kind Code	B1
Date of Patent	September 02, 2025
Inventor(s)	Shah; Samir et al.

Decision tree training using a database system

Abstract

In an embodiment, a computer-implemented method for training a decision tree using a database system, the decision tree comprising a plurality nodes, comprises, by one or more computing devices: storing in a database input data for training the decision tree, the input data comprising a plurality of feature values corresponding to a plurality of features; generating a particular node of the plurality of decision nodes by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data; identifying a particular feature of the subset of the plurality of features associated with the highest variance value; associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature.

Inventors: Shah; Samir (San Francisco, CA), Zhang; Bei (San Jose, CA), Miller; Kenton (Sunnyvale, CA)

Applicant: Shape Security, Inc. (Mountain View, CA)

Family ID: 96882204

Assignee: Shape Security, Inc. (Seattle, WA)

Appl. No.: 16/457884

Filed: June 28, 2019

Publication Classification

Int. Cl.: G06N20/20 (20190101); G06N5/01 (20230101); G06N20/00 (20190101)

U.S. Cl.:

CPC G06N20/20 (20190101); G06N5/01 (20230101); G06N20/00 (20190101);

Field of Classification Search

CPC: G06N (20/20); G06N (5/003)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6108648	12/1999	Lakshmi et al.	N/A	N/A
10733320	12/2019	Nerurkar	N/A	N/A
11061902	12/2020	Idicula	N/A	N/A
2002/0120435	12/2001	Frazier	N/A	N/A
2006/0003419	12/2005	Pan	N/A	N/A
2006/0034194	12/2005	Kahan	N/A	N/A
2006/0230006	12/2005	Buscema	N/A	N/A
2007/0017996	12/2006	Xia	N/A	N/A
2007/0179966	12/2006	Li et al.	N/A	N/A
2009/0281981	12/2008	Chen et al.	N/A	N/A
2014/0172753	12/2013	Nowozin	N/A	N/A
2015/0278635	12/2014	Poggio et al.	N/A	N/A
2015/0310044	12/2014	Isaacson et al.	N/A	N/A
2017/0022822	12/2016	Taffet et al.	N/A	N/A
2017/0126627	12/2016	Yang	N/A	H04L 63/20
2017/0228224	12/2016	Dunn et al.	N/A	N/A
2018/0032870	12/2017	Liu	N/A	N/A
2018/0082172	12/2017	Patel et al.	N/A	N/A
2018/0306609	12/2017	Agarwal	N/A	G06N 20/10
2019/0026489	12/2018	Nerurkar	N/A	N/A
2019/0057356	12/2018	Larsen	N/A	G06V 40/28
2019/0188000	12/2018	Chen	N/A	N/A
2019/0318104	12/2018	Yoshino	N/A	G06N 20/20
2019/0347327	12/2018	Patil	N/A	G06F 16/951
2020/0073953	12/2019	Kulkarni	N/A	G06N 3/084

OTHER PUBLICATIONS

Sugumaran, V., and K. I. Ramachandran. "Effect of No. of features on classification of roller bearing faults using SVM and PSVM." Expert Systems with Applications 38.4 (2011): 4088-4096. (Year: 2011). cited by examiner

Bentayeb, Fadila, and Jérôme Darmont. "Decision tree modeling with relational views." International Symposium on Methodologies for Intelligent Systems. Springer, Berlin, Heidelberg, 2002. (Year: 2002). cited by examiner

Chen, Engu, and Xin James He. "Crude oil price prediction with decision tree based regression approach." Journal of International Technology and Information Management 27.4 (2019): 2-16. (Year: 2019). cited by examiner

Sugumaran, V., and K. I. Ramachandran. "Effect of number of features on classification of roller bearing faults using SVM and PSVM." Expert Systems with Applications 38.4 (2011): 4088-4096. (Year: 2011). cited by examiner

Doyle, Michelle, and Pdraig Cunningham. "A dynamic approach to reducing dialog in on-line

decision guides.” European Workshop on Advances in Case-Based Reasoning. Springer, Berlin, Heidelberg, 2000. (Year: 2000). cited by examiner

McClean, Sally I. “Data mining and knowledge discovery.” (2003): 229-246. (Year: 2003). cited by examiner

Zaklouta, Fatin, and Bogdan Stanciulescu. “Real-time traffic-sign recognition using tree classifiers.” IEEE Transactions on Intelligent Transportation Systems 13.4 (2012): 1507-1514. (Year: 2012). cited by examiner

Basak, Jayanta, and Raghu Krishnapuram. “Interpretable hierarchical clustering by constructing an unsupervised decision tree.” IEEE transactions on knowledge and data engineering 17.1 (2005): 121-132. (Year: 2005). cited by examiner

Von Hacht, August, “Implementation of Deep Feedforward Neural Network with CUDA Backend for Efficient and Accurate Natural Image Classification,” Department of Physics, Chalmers University of Technology, 69 pp, Jun. 2017. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/064552, filed Dec. 4, 2019, mailed Mar. 17, 2020. cited by applicant

International Preliminary Report on Patentability for PCT/US2019/064552, dated m Jun. 2021. cited by applicant

Dabbura, Imad, “Coding Neural Network—Forward Propagation and Backpropagation,” <https://towardsdatascience.com/coding-neural-network-forward-propagation-and-backpropagation-ccf8cf369f76>, 19 pp, Mar. 31, 2018. cited by applicant

International Search Report and Written Opinion for International Application No. PCT/US2019/064550, filed Dec. 4, 2019, mailed Feb. 12, 2020. cited by applicant

Von Hacht, August, “Implementation of Deep Feedforward Neural Network with CUDA Backend for Efficient and Accurate Natural Image Classification,” Department of Physics, Chalmers University of Technology, 69 pp, Jun. 2017. cited by applicant

Search Report and the Written Opinion of corresponding PCT Application Serial No. PCT /US 19/64552, mailed Mar. 17, 2020. cited by applicant

Imani, Mohsen et al. “RAPI DN N: In-memory Deep Neural Network Acceleration Framework” arXiv [Version 3 published 2018] [Retrieved Apr. 2022] <URL:<https://arxiv.org/abs/1806.05794>> (Year: 2018). cited by applicant

Haridas, Harisanker “Deep Neural Network implemented in pure SQL over BigQuery”. Towards DataScience [Published Mar. 2018] [Retrieved Apr. 2022] <URL: <https://towardsdatascience.com/deep-neural-network-implemented-in-pure-sql-over-bigquery-f3ed245814d3>> (Year: 2018). cited by applicant

Carlile, Brad et al. “Columnar Database Techniques for Creating AI Features” arXiv [Published 2017] [Retrieved Apr. 2022] <URL: <https://arxiv.org/abs/1712.02882>> (Year: 2017). cited by applicant

M. S. Razlighi, M. Imani, F. Koushanfar and T. Rosing, “LookNN: Neural network with no multiplication,” Design, Automation & Test in Europe Conference & Exhibition (DATE), 2017, 2017, pp. 1775-1780, doi: 10.23919/DATE.2017.7927280. (Year: 2017). cited by applicant

Sergio Fernandes and Jorge Bernardino. 2015. What is BigQuery? In <i>Proceedings of the 19th International Database Engineering Applications Symposium</i> (<i>IDEAS '15</i>}. Association for Computing Machinery, New York, NY, USA, 202-203. DOI :<https://doi.org/10.1145/2790755.2790797> (Year: 2015). cited by applicant

Bagherinezhad, Hessam et al. “LCNN: Lookup-based Convolutional Neural Network” arXiv [Published 2017] [Retrieved Apr. 2022] <URL: <https://arxiv.org/abs/1611.06473>> (Year: 2017). cited by applicant

SQL Server Team, “How six lines of code+ SQL Server can bring Deep Learning to ANY App”, 2017, Microsoft, retrieved from <https://www.microsoft.com/en-us/sql-server/blog/2017/01/05/how-six-lines-of-code-sql-server-can-bring-deep-learning-to-any-app/> on Jun. 20, 2024 (Year: 2017).

cited by applicant

Canali, Luca, "How to Build a Neural Network Scoring Engine in PUSQL", 2016, retrieved from <https://externaltable.blogspot.com/2016/07/a-neural-network-scoring-engine-in-plsql.html> on Jun. 20, 2024 (Year: 2016). cited by applicant

Shevchuk, Yuriy, "Neural Networks in MySQL", 2016, retrieved from <http://blog.itdxxr.com/2016/07/01/neural-networks-in-mysql.html> on Jun. 20, 2024 (Year: 2016). cited by applicant

Chen, "Design of Resistive Synaptic Devices and Array Architectures for Neuromorphic Computing", May 2018, Arizona State University (Year: 2018). cited by applicant

Kim et al., "Input-Splitting of Large Neural Networks for Power-Efficient Accelerator with Resistive Crossbar Memory Array", Jul. 23, 2018, ISLPED '18: Proceedings of the International Symposium on Low Power Electronics and Design, vol. 2018, pp. 1-6 (Year: 2018). cited by applicant

Gao et al., "Fully parallel write/read in resistive synaptic array for accelerating on-chip learning", 2015, Nanotechnology, vol. 26, pp. 1-9 (Year: 2015). cited by applicant

Primary Examiner: Huntley; Michael J

Assistant Examiner: Kim; Sehwan

Attorney, Agent or Firm: Troutman Pepper Locke LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) This application is related to U.S. application Ser. No. 16/222,974, filed Dec. 17, 2018 titled "Decision Tree Training Using a Database System," the entire contents of which is hereby incorporated by reference as if fully set forth herein.

FIELD OF THE DISCLOSURE

(2) One technical field of the disclosure is computer-implemented machine learning systems such as neural networks and classifiers. Another technical field is computer-implemented methods of training a decision tree model. Still another field is databases, such as relational databases.

BACKGROUND

(3) The approaches described in this section are approaches that could be pursued, but not necessarily approaches that have been previously conceived or pursued. Therefore, unless otherwise indicated, it should not be assumed that any of the approaches described in this section qualify as prior art merely by virtue of their inclusion in this section.

(4) Machine learning techniques provide a solution to the difficulties of solving complex problems. Machine learning may be used in various kinds of data analysis systems, voice recognition, document recognition, image recognition, recommendations systems, and other applications of artificial intelligence. One technique utilizes decision trees or random forests, which comprise a plurality of decision trees. A decision tree generally contains a plurality of nodes and branches, where each node tests an attribute of input data and branches to other nodes based on the outcome of the test.

(5) A benefit of decision trees is that they learn inherent patterns in different datasets. Thus, a programmer does not have to know the correct attributes and/or branching methods prior to programming the decision tree for the decision tree to provide an accurate result.

(6) However, while machine learning techniques such as decision trees provide a large benefit in

solving complex problems, they do so at a high computational cost. Training the decision tree(s) requires a large amount of computing power and memory. The calculations required take a significant amount of both time and processing power to compute. For example, large numbers of training examples are typically used to train decision trees in order to carry out classification or regression tasks. Each training example may include a large number of features, and calculations may be performed for each of the features. Thus, the training process is typically time consuming and resource intensive.

(7) Therefore, there is a need for a technique that reduces the computational cost of training decision trees.

SUMMARY

(8) The appended claims may serve as a summary of the disclosure.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) In the drawings:

(2) FIG. 1 is a block diagram illustrating an example decision tree, according to an embodiment.

(3) FIG. 2 is a block diagram illustrating an example random forest, according to an embodiment.

(4) FIG. 3 is a block diagram illustrating an example system in which the techniques described herein may be practiced, according to an embodiment.

(5) FIG. 4 is a block diagram illustrating an example decision tree training process, according to an embodiment.

(6) FIG. 5 is a block diagram that illustrates a computer system upon which an embodiment of the invention may be implemented.

DETAILED DESCRIPTION

(7) In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present invention. It will be apparent, however, that the present invention may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form in order to avoid unnecessarily obscuring the present invention. Embodiments are disclosed in sections according to the following outline: 1. GENERAL OVERVIEW 2. DECISION TREE OVERVIEW 3.

DECISION TREE TRAINING USING VARIANCE 4. OVERVIEW OF EXAMPLE

EMBODIMENTS 4.1 DECISION TREE TRAINING MODULE 4.2 DATA ENGINE 4.3 DATA STORE 5. PROCESS OVERVIEW 6. HARDWARE OVERVIEW

1. General Overview

(8) Techniques are described herein for training one or more decision trees using a database system. In an embodiment, input data for training the one or more decision trees is stored in a database, the input data comprising a plurality of feature values corresponding to a plurality of features. A particular decision node is generated by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data; identifying a particular feature of the subset of the plurality of features associated with the highest variance value; and associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature.

(9) In an embodiment, calculating a variance value associated with the feature comprises partitioning the subset of the input data based on the feature. An aggregate label is applied to each partition of the subset, and the variance value is based on the output of the aggregate labels. An aggregate label is a function that is defined on collection of records. In some embodiments, the aggregate label may be one of: Fourier mode, login success, entropy quantity, or manual fraud

signature.

(10) In an embodiment, the selecting, calculating, identifying, and associating are performed for each decision node of a plurality of decision nodes of a decision tree. Additionally, the selecting, calculating, identifying, and associating may be performed in parallel for multiple decision nodes.

(11) In an embodiment, after generating the plurality of decision nodes of the decision tree, one or more queries and/or one or more program instructions corresponding to the decision tree are generated. The one or more queries and/or one or more program instructions corresponding to the decision tree, when executed, applies the decision tree to a set of input data.

(12) In an embodiment, to generate a random forest, the selecting, calculating, identifying, and associating are performed for each decision node of a plurality of decision nodes of a plurality of decision tree. The selecting, calculating, identifying, and associating may be performed in parallel for multiple decision nodes and/or multiple decision trees.

(13) Additionally, after generating the plurality of decision trees, one or more queries and/or one or more program instructions corresponding to the random forest are generated. The one or more queries and/or one or more program instructions corresponding to the random forest, when executed, applies the plurality of decision trees to a set of input data.

2. Decision Tree Overview

(14) A decision tree is a machine learning technique that utilizes a set of decision nodes. The decision nodes are arranged into a plurality of layers in a tree, or flow chart, structure. A root node is connected to a plurality of child nodes in a hierarchical structure, with each decision node branching into two or more child nodes, and ending at a plurality of leaf nodes.

(15) Each decision node receives input data and tests a particular attribute of the input data. The decision node may generate and/or store output corresponding to the outcome of the test. Additionally, if the decision node is not a leaf node, a decision is made as to which child node of the decision node to branch to based on the outcome of the test.

(16) FIG. 1 illustrates an example decision tree. In the illustrated example, decision tree **100** comprises seven nodes, node **102-114**, arranged in three layers, layer **122**, layer **124**, and layer **126**. Each of nodes **102-114** may be configured to test a particular attribute, feature, or value of input data received at the node, and generate output data corresponding to the outcome of the test. Additionally, non-leaf nodes such as nodes **102-106** may branch to a node in the subsequent layer based on the outcome of the test. In other embodiments, more or fewer layers may be used, and the number of nodes and/or branches at each layer may differ based on the embodiment.

(17) In FIG. 1, node **102** receives input **130**. In an embodiment, input **130** comprises a plurality of input values. For example, input **130** may be a vector, array, database table, or other multi-dimensional object. Node **102** may be configured to test a particular attribute, feature, or value of the plurality of input values. Based on the outcome of the test, one of child nodes **104** and **106** is selected as the next node to traverse, and input **130** is provided to the selected node. Additionally, output may be generated corresponding to node **102**.

(18) For the purpose of illustrating a clear example, assume input **130** comprises three values corresponding to three features: x, y, and z. Node **102** may be configured to determine whether the value of “x” is greater than or less than 3. If “x” is less than 3, then node **102** may branch to node **104**, and if “x” is greater than or equal to 3, then node **102** may branch to node **106**. Each of node **104** and **106** may branch based on different tests and/or different attributes, features, or values. For example, node **104** may branch to either node **108** or node **110** based on whether the value of “y” is greater than or less than 5, and node **106** may branch to either node **112** or node **114** based on whether the value of “y” is greater than or less than 10.

(19) After reaching a leaf node, such as nodes **108-114**, output of the decision tree is generated based on the decision nodes that were traversed or selected. In an embodiment, the output of the decision tree comprises the set of output generated by the selected decision nodes. In other embodiments, the output of the decision tree is generated based on the output of the selected

decision nodes. For example, the output may be generated by determining a majority-vote of the output of the traversed decision nodes or by averaging the output of the traversed decision nodes. (20) As an example, assume decision tree **100** is trained to determine a yes or no categorization based on input **130**. Each of nodes **102-114** may generate output indicating a respective yes or no determination. The output **140** of the decision tree **100** may be based on whether more selected nodes generated a yes determination versus a no determination. Alternately, each node may generate an output of 0 corresponding to a no determination and output of 1 corresponding to a yes determination. The output **140** of the decision tree **100** may be an average of the output generated by selected decision nodes, and indicate a percentage likelihood of a yes categorization. For example, assume node **102** generated an output of 0, node **104** generated an output of 1, and node **110** generated an output of 0. The output **140** of the decision tree **100** may be the average of the three output values, or 0.33.

(21) A random forest is a machine learning technique that utilizes a plurality of decision trees. Input data is provided to each decision tree of the plurality of decision trees, and each decision tree trains a respective output. The output of the random forest is generated based on the plurality of decision trees. In an embodiment, the output of the random forest comprises the set of output generated by the plurality of decision trees. In other embodiments, the output of the random forest is generated based on the output of the plurality of decision trees. For example, the output may be generated by determining a majority-vote of the output of the decision trees or by averaging the output of the decision trees.

(22) FIG. 2 illustrates an example random forest **200**. In the illustrated example, random forest **200** comprises three decision trees, decision tree **202**, decision tree **204**, and decision tree **206**. Each of decision tree **202**, decision tree **204**, and decision tree **206** comprises a set of decision nodes, respectively. In other embodiments, random forest **200** may comprise more or fewer decision trees, and each decision tree may comprise more or fewer nodes. Additionally, the number of nodes and/or layers at each decision tree may be different.

(23) Each of decision tree **202**, decision tree **204**, and decision tree **206** receives input **220** and generates output, respectively. Output **230** of the random forest is generated based on outputs.

(24) As an example, assume random forest **200** is configured to determine a yes or no categorization based on input **220**. Each of output may indicate a respective yes or no determination. The output **230** of the random forest **200** may be based on whether more decision trees generated a yes determination versus a no determination. For example, if output **214** indicates a yes determination and outputs indicate a no determination, then output **230** may indicate a no determination. Alternately, the output **230** of the random forest **200** may be an average of the outputs. For example, each output may be a value of 0 for a no determination and a value of 1 for a yes determination. Assume output **214** is 0, output **212** is 1, and output **216** is also 0. The output **230** of the random forest **200** may be the average of the three decision tree output values, or 0.33.

(25) Each decision tree of the random forest may branch in different ways. That is, while each decision tree may receive the same set of input data, the root node of each decision tree may test a different attribute, feature, or value of the input data. Similarly, subsequent decision nodes may also test a different attribute, feature or value of the input data compared to decision nodes of the same level and/or location in the other decision tree(s).

3. Decision Tree Training Using Variance

(26) Training a machine learning model comprises providing the machine learning model with a set of training data. The training data may comprise a plurality of input data and a plurality of corresponding output data. The machine learning model may use the training inputs and the training outputs to infer a method for correlating the input to the output. Training a decision tree comprises determining an attribute or feature of the input data on which each decision node should branch.

(27) Typically, to train a decision tree, information gain is calculated for each feature of a plurality

of features of a set of training input data based on corresponding training output data. This is referred to as supervised training, where training is based on labels that indicate, for each item of training input data, an output associated with the item of training input data, e.g. a category, true/false value, etc. Thus, generating the training output data requires determining a label or other output value and associating the label, or other output value, with each item of training input data. Furthermore, the training output data must be stored alongside the training input data.

(28) In an embodiment, decision tree training is performed unsupervised or semi-supervised. Unsupervised and/or semi-supervised training does not require each item of input data to be associated with a corresponding output. Instead, the decision tree is trained based on identifying patterns and/or anomalies in the training input data. Thus, in some embodiments, the set of training data may include only training input data and may not include training output data.

(29) In an embodiment, unsupervised training of a decision tree comprises applying an aggregate label to the set of input data, or a portion of the subset of input data. An aggregate label is a function that is defined on collection of records. As an example, an aggregate label may be, for example, an average, mean, minimum, maximum, etc. of a set of values. As another example, an aggregate label may describe how or whether the set of values conform to a particular pattern. Applying an aggregate label to a set of data generates an aggregate output for the set of data, i.e. a label describing the set of data.

(30) In an embodiment, training a decision tree comprises calculating a variance for each feature of a plurality of features of the set of training input data. Variance for a particular feature indicates, when the set of training input data is divided based on the particular feature, how much the divided portions of training input data differ from each other.

(31) In an embodiment, to calculate variance of a feature, the training input data is divided based on the feature, and an aggregate label is applied to each divided portion of the training input data.

(32) For example, assume the condition “contains ‘foo’” is introduced for a feature Y. The input data may be divided into two subsets: a first subset where the value of the feature Y contains ‘foo,’ and a second subset where the value of the feature Y does not contain ‘foo.’ An aggregate label is applied to the first subset to generate a first aggregate label output value, and to the second subset to generate a second aggregate label output value. Variance of feature Y is calculated based on the first aggregate label output value and the second aggregate label output value.

(33) In an embodiment, calculating the variance comprises dividing the first aggregate label output value by the number of elements in the first subset, dividing the second aggregate label output value by the number of elements in the second subset, and calculating the difference between the two results.

(34) Referring to the above example, assume an aggregate label A finds the average of a set of data. Each value of a feature Y defines a binary split of the data set D into two subsets, D1 and D2. Considering the aggregate label A as a random variable over {D1, D2}, the variance of A is calculated for all values of the feature Y. An example equation for calculating the average of A may be:

$$\text{Avg}(A)=p1*a1+p2*a2$$

where p1 is the probability of D1 and p2 is the probability of D2, and aggregate label A takes the value of a1 on D1 and the value of a2 on D2. An example equation for calculating the variance of A based on the average of A may be:

$$\text{Var}(A)=\text{Avg}((A-\text{Avg}(A)).\text{sup.2})=p1*(a1-\text{Avg}(A)).\text{sup.2}+p2*(a2-\text{Avg}(A)).\text{sup.2}$$

(35) In an embodiment, the decision tree is trained to receive input comprising request data corresponding to a request from a client computer to a server computer, and determine whether the request is legitimate or illegitimate based on the request data. As an example, a legitimate request may be from browser of a legitimate user, while an illegitimate request may be from a headless browser or a bot. The aggregate label(s) used to train the decision tree may be based on characteristics of requests to the server computer. Additionally or alternatively, the aggregate

label(s) may identify patterns in a set of requests to the server computer. Example aggregate labels based on requests to a server computer include, but are not limited to, Fourier mode, login success, entropy quantity, and manual fraud signatures. Any number and/or types of aggregate labels may be used, depending on the implementation, and are not 1.

(36) A login success aggregate label determines what fraction of a set of training input data included successful login attempts. For example, if the training input data comprises request data corresponding to 100 login requests, and 80 requests succeeded in logging in, then the login success aggregate label may generate output indicating a percentage of 80%.

(37) A Fourier mode aggregate label determines how well a set of training input data conforms to a diurnal pattern, defined as projection of the traffic onto the diurnal Fourier mode. A diurnal pattern refers to a pattern that reoccurs over a period of time, i.e. a cyclical pattern. For example, the amount of network traffic might peak during certain times of the day and ebb during other times of the day. The network traffic over a period of time may be visualized as a smooth wave with peaks and ebbs. In contrast, a non-diurnal pattern may have intermittent spikes or other anomalies, and does not conform to any cyclical pattern. As an example, unusual (e.g. illegitimate) network traffic may spike at odd times. The non-conforming network traffic may be visualized as sharp spikes or plateaus within the (smooth) wave at particular times. In an embodiment, training decision trees with the Fourier diurnal mode comprises partitioning a high-dimensional time series dataset by feature values according to their time series pattern.

(38) An entropy quantity aggregate label determines an entropy of mouse movement characteristics of a set of training input data. Mouse movement characteristics describes mouse movements at a client computer prior to sending the request. For example, the entropy quantity aggregate label may generate output indicating high entropy if the requests in the training input data included different mouse movements, while the entropy quantity aggregate label may generate output indicating low entropy if similar (or no) mouse movement occurred in several requests of the training input data.

(39) A manual fraud signatures aggregate label determines a fraction of keyboard events that include one or more sets of particular keystrokes. For example, the manual fraud signatures aggregate label may determine what fraction of requests in the training input data included “ctrl+c” and “ctrl+v” (copy and paste keystrokes).

(40) In an embodiment, training a decision tree comprises calculating the variance associated with each feature of a plurality of features of the input data, and determining a particular feature of the plurality of features that resulted in the highest variance. The feature of the input data with the highest variance is selected as a root node. The root node branches based on the selected feature.

(41) Additionally, calculating the variance may comprise identifying a plurality of conditions for splitting the input data based on each feature of the plurality of features, and identifying the condition associated with the highest variance. The condition may be a particular value, or a particular range of values, for the feature, depending on the data type associated with the feature. For example, if a feature is associated with numeric data, the condition may be a range. If a feature is associated with string data, the condition may be whether the string includes a particular string value.

(42) Referring to the above example, variance is calculated for each binary split of the feature Y. For example, assume the feature Y is categorical and takes on three possible values: “foo,” “bar,” and “car.” The condition $Y = \text{“foo”}$ defines a binary split of the training set D into D1 and D2, where D1 comprises data whose value of feature Y is “foo,” and D2 comprises data whose value of feature Y is not “foo.” Variance is calculated based on the sets D1 and D2 for $Y = \text{“foo”}$. This is repeated for $Y = \text{“bar”}$ and $Y = \text{“car”}$. The split that results in the highest variance determines how the node branches. For example, if $Y = \text{“bar”}$ results in the highest variance, then $Y = \text{“bar”}$ defines the tree split: one branch corresponds to when the value of feature Y is “bar,” and the other branch corresponds to when the value of feature Y is not “bar.”

(43) The input data is divided based on the branches of the root node. For each branch, variance is

calculated for each of the remaining features based on the portion of the divided input data corresponding to the branch. The feature of the input data with the highest variance is selected as the decision node for the branch. For example, assume the root node branches based on the condition “ $Y < 3$ ”. The input data is divided into two subsets: a first subset where the value of feature Y is less than 3, and a second subset where the value of feature Y is greater than or equal to 3. A decision node is generated based on the first subset, and a second decision node is generated based on the second subset.

(44) In some embodiments, the same aggregate label is used to calculate variance for each decision node. For example, a Fourier mode aggregate label may be used when training all decision nodes of the decision tree.

(45) In other embodiments, a plurality of aggregate labels may be available, and a particular aggregate label is selected for use in calculating variance at each decision node. For example, a Fourier mode aggregate label may be used for training the root node, a login success aggregate label may be used for training the first child node of the root node, and an entropy quantity aggregate label may be used for training the second child node of the root node. Any technique may be used for selecting the aggregate label used. For example, for each decision node, a particular aggregate label may be selected randomly.

(46) The process of dividing the input data into subsets and generating a decision node corresponding to each subset is repeated for each branch of each decision node, until no features remain or the variance of all remaining features is 0.

(47) Additionally or alternatively, rather than generating decision nodes by following the generated branches and dividing the input data, each decision node may be trained independently from other decision nodes in the tree. In an embodiment, each decision node is generated based on a subset of features of the input data. Additionally or alternatively, each decision node may be generated based on a subset of the input data. Selecting the subset of features and/or the subset of input data may be performed for each decision node, and the set of features and/or the subset of input data may be different for each decision node. Additionally, the subset of features and/or the subset of input data may be selected randomly for each decision node. For each decision node, variance is calculated for each feature of the input data with the highest variance is selected as the feature for that decision node.

(48) In an embodiment, a number of decision nodes for the decision tree may be determined and the steps described above are performed a number of times based on the determined number of decision nodes. For example, for decision tree **100**, since it comprises three levels and each node branches in two, decision tree **100** comprises seven decision nodes. The steps of selecting a subset of features and/or input data, applying an aggregate label, calculating variance, and determining a feature corresponding to the highest variance are performed seven times to generate the seven decision nodes. After the decision nodes are generated, the decision nodes may be connected based on their location in the decision tree.

(49) In an embodiment, to train a random forest, each decision tree of the random forest is trained separately. The generated decision trees are combined to form the random forest. For example, referring to random forest **200**, each of decision tree **202**, decision tree **204**, and decision tree **206** may be generated separately. The steps described above are performed for each of decision tree. Additionally, the steps may be performed in parallel for each decision tree.

(50) In an embodiment, training input data comprises a plurality of attributes or features and each decision tree of the random forest is trained based on a different subset of the plurality of attributes or features. For example, assume input data **220** comprises ten features, A-J. Decision tree **202** may be trained to test features A-G, decision tree **204** may be trained to test features A-D and F-H, and decision tree **206** may be trained to test features D-J.

4. Overview of Example Embodiments

(51) In an embodiment, a system **300** comprises components that are implemented at least partially

by hardware at one or more computing devices, such as one or more hardware processors executing stored program instructions stored in one or more memories for performing the functions that are described herein. In other words, all functions described herein are intended to indicate operations that are performed using programming in a special-purpose computer or general-purpose computer, in various embodiments. FIG. 3 illustrates only one of many possible arrangements of components configured to execute the programming described herein. Other arrangements may include fewer or different components, and the division of work between the components may vary depending on the arrangement.

(52) In FIG. 3, system **300** comprises decision tree training module **310**, random forest **320**, data engine **330**, data store **340**, and application program **350**. In an embodiment, each of decision tree training module **310**, random forest **320**, data engine **330**, data store **340**, and application program **350** are located at or executing on one or more computing devices. Additionally or alternatively, two or more of decision tree training module **310**, random forest **320**, data engine **330**, data store **340**, and application program **350** may be located at or executing on the same or overlapping sets of one or more computing devices. Example computing devices include, but are not limited to, servers, racks, workstations, personal computers, laptops, wired devices, wireless devices such as mobile phones or tablets, multi-processor systems, and the like. Additionally or alternatively, one or more of decision tree training module **310**, random forest **320**, data engine **330**, data store **340**, and application program **350** may be located at or executing on one or more virtual computing instances, such as Microsoft Azure, Oracle Cloud, Amazon Elastic Compute Cloud (EC2), or Google Compute Engine instances.

4.1 Decision Tree Training Module

(53) In an embodiment, a decision tree training module **310** comprises program instructions that are programmed or configured to receive training input data and generate one or more decision trees. Additionally or alternatively, decision tree training module **310** comprises program instructions that are programmed or configured to receive training input data and generate one or more random forests, each random forest comprising a plurality of decision trees.

(54) In the embodiment illustrated in FIG. 3, decision tree training module **310** receives training data, such as input data **342**, and generates a random forest **320**. In FIG. 3, random forest **320** comprises decision tree **322**, decision tree **324**, and decision tree **326**. Each of decision trees **322**, **324**, and **326** comprise a respective plurality of nodes **328**. Each of the nodes **328** may be associated with a particular feature of input data **342** and branch based on applying a condition to the particular feature of the input data **342**.

(55) In an embodiment, decision tree training module **310** is programmed or configured to generate a decision tree based on a subset of the input data **342**. Additionally or alternatively, decision tree training module **310** may be programmed or configured to generate the decision tree based on a subset of a plurality of features of the input data **342**. For example, assume input data **342** comprises 1000 sets of input, each input comprising 50 features. Decision tree training module **310** may generate decision tree based on 20 features and 500 sets of input from the input data. The subset of features and/or the subset of the input data may be selected randomly.

(56) In an embodiment, decision tree training module **310** is programmed or configured to generate the layers of a decision tree in sequential order. Generating decision nodes of each layer may be based on dividing the training data based on how the previous decision nodes branch, as discussed above. If a subset of features and/or a subset of training data are used to generate the decision tree, the same subset(s) may be used to generate all nodes of the decision tree. If decision tree training module **310** generates a random forest, the subset(s) may be selected separately for each decision tree.

(57) In other embodiments, decision tree training module **310** is programmed or configured to train nodes of a decision tree separately. If a subset of features and/or a subset of training data are used to generate the decision tree, the subset(s) may be selected separately for each decision node. Each

decision node may therefore be generated based on a different subset of features and/or training data.

(58) Additionally, the subset of training data selected for each node may be filtered based on how one or more previous decision nodes branch. For example, a parent node may branch based on whether the value of feature A is greater than or less than 3. A first child node may be trained using a subset of training data where A is greater than or equal to 3. A second child node may be trained using a subset of training data where A is less than 3. However, the first child node and the second child node may each be trained using a different subset of features. For example, the first child node may be trained based on features C, D, and E, while the second child node may be trained based on features D, F, and G.

(59) Since training each branch of the decision tree is not based on what features are selected at each other branch, some or all of the branches of the decision tree may be generated in parallel. Referring to the above example, after the parent node is trained, the first and second child nodes may be trained in parallel. If decision tree training module **310** generates a random forest, some or all of the decision trees of the random forest may be generated in parallel. Thus, at any given time, multiple nodes may be training, both across decision trees and within a particular decision tree.

(60) In an embodiment, decision tree training module **310** comprises one or more queries **312** which, when executed, performs one or more functions for training a decision tree, such as: retrieving training input data **342** from data store **340**, applying an aggregate label to each feature of one or more features of input data **342**, calculating variance for each feature of the one or more features of input data **342**, identifying a feature associated with the highest variance, etc.

(61) For the purpose of illustrating a clear example, assume a feature is associated with a set of discrete feature values. A query **312** may, when executed by system **300**, identify the set of discrete feature values and, for each discrete feature value, divide a set of input data based on the discrete feature value, apply an aggregate label to each half of the divided set of input data, and calculate a variance associated with the discrete feature value.

(62) In an embodiment, one or more queries may each be configured to calculate variance based on a particular aggregate label. In FIG. 3, the decision tree training queries **312** comprise a login success aggregate label query **312A**, Fourier mode aggregate label query **312B**, manual fraud signatures aggregate label query **312C**, and entropy quantity aggregate label query **312D**. Each query uses a respective aggregate label—login success aggregate label query **312A** uses a login success aggregate label, Fourier mode aggregate label query **312B** uses a Fourier mode aggregate label, manual fraud signatures aggregate label query **312C** uses a manual fraud signatures aggregate label, and entropy quantity aggregate label query **312D** uses an entropy quantity aggregate label.

(63) Each query may be configured to divide a set of input data based on a particular feature, apply a respective aggregate label to each divided set of input data, and calculate variance based on applying the respective aggregate label to the divided sets. Additionally, each query may be configured to identify one or more candidate feature values (for branching a decision node), and calculate the variance for the one or more candidate features. The method by which the one or more candidate feature values are identified may vary depending on the associated data type of the particular feature and the particular implementation.

(64) In an embodiment, for a numeric data type, identifying the one or more candidate feature values comprises sorting the training input data based on the particular feature and selecting feature values at regular intervals from the sorted input data. In other embodiments, other methods for identifying a plurality of values may be used, such as identifying frequently occurring or common values, an average value, minimum and maximum values, etc. The query may be configured to identify a plurality of candidate values on which the training input data may be split, calculate a variance associated with each candidate value of the plurality of candidate values, and generate output indicating the plurality of candidate values and their associated variance. The highest

variance may be used as the variance of the particular feature, when comparing against the variance of other features. If the particular feature is selected as a decision node, the decision node may branch based on the candidate value that was associated with the variance. For example, the decision node may branch based on whether a feature value of the particular feature is greater, equal to, and/or less than the candidate value.

(65) In an embodiment, for a categorical data type, identifying the one or more candidate feature values comprises identifying a plurality of discrete values for the particular feature in the training input data. The query may be configured to identify a plurality of discrete values for the particular feature in the training input data, calculate a variance associated with each discrete value, and generate output indicating the plurality of discrete values and their associated variance. The highest variance may be used as the variance of the particular feature, when comparing against the variance of other features. If the particular feature is selected as a decision node, the decision node may branch based on the discrete value that was associated with the highest variance. For example, the decision node may branch based on whether a feature value of the particular feature is the discrete value.

(66) In an embodiment, for a string data type, identifying the one or more candidate features values comprises identifying a plurality of string tokens from the feature values of the particular feature in the training input data. The query may be configured to identify a plurality of string tokens from the feature values of the particular feature in the training input data, calculate variance associated with each string token, and generate output indicating the plurality of string tokens and their associated variance. The highest variance may be used as the variance of the particular feature, when comparing against the variance of other features. If the particular feature is selected as a decision node, the decision node may branch based on the string token that was associated with the highest variance. For example, the decision node may branch based on whether a feature value of the particular feature includes the string token.

(67) In an embodiment, identifying the plurality of tokens comprises determining a plurality of string values for the particular feature from the training input data. Each string value may be split into one or more tokens. String values may be split on, for example, spaces; delimiters such as commas, periods, semi-colons, and dashes; punctuation marks such as question marks, periods, and exclamation points; special characters or symbols such as tabs and carriage returns; or any other characters or expressions. The plurality of tokens may comprise one or more unique tokens from the split string values.

(68) In an embodiment, each query may each be configured to calculate variance for a particular data type. Each query may be configured to process data of the respective data type. For example, a first query and a second query may both be configured to apply a Fourier mode aggregate label. The first query may be configured to process numerical data types and the second query may be configured to process string data types.

(69) Alternately, each query may be configured to calculate variance for a plurality of data types. The query may be configured to receive input specifying a particular feature, determine a data type associated with the particular feature, and identify one or more candidate feature values for the particular feature based on the data type associated with the particular feature.

(70) In an embodiment, the decision tree training module **310** is programmed or configured to generate one or more program instructions which, when executed, applies the generated decision tree(s) and/or random forest(s) to input data. Additionally or alternatively, the decision tree training module **310** may be programmed or configured to generate one or more queries which, when executed by a database system, applies the generated decision tree(s) and/or random forest(s) to input data. In some embodiments, the decision tree training module **310** generates the one or more queries, and the one or more queries may be transformed by the decision tree training module **310** or other applications, programs, or modules, into one or more program instructions which, when executed, applies the generated decision tree(s) and/or random forest(s) to input data.

- (71) For the purpose of illustrating a clear example, assume a decision node tests whether a particular feature value contains the string token “537.” An example portion of a query corresponding to the decision node may be: IF(IFNULL(pv.nr.ut,“default value”) LIKE “%537%”, 1.0, 0.0)
- (72) In the above example, assume “pv.nr.ut” refers to the column corresponding to the particular feature. The IF function determines whether the value of the input data in column “pv.nr.ut” contains the token “537.” If the value contains the specified token, then the IF function returns an output of 1.0. If the value does not contain the specified token, then the IF function returns an output of 0.0. Additionally, if the feature “pv.nr.ut” is null—that is, the feature does not exist in the input data, then a default value is used.
- (73) In an embodiment, a query corresponding to a decision tree comprises a plurality of nested IF functions. Additionally, a query corresponding to a random forest may comprise a plurality of sub-queries, each sub-query corresponding to a particular decision tree. The output of the plurality of sub-queries may be averaged to generate output of the random forest.
- (74) An example nested IF function may be expressed in one example as:
IF(foo.parameters.params.vndr LIKE % default %”, IF(pv.nr.ut LIKE “%537%”, 1.0, 0.0)
IF(fooElement.barHash=77957379.0, 1.0, 0.0))
- (75) In the above example, a first IF function determines whether the value of the input data in the column “foo.parameters.params.vndr” contains a default value string token. If the value contains the specified token, then the first IF function proceeds to a second IF function. As discussed above, the second IF function determines whether the value of the input data in column “pv.nr.ut” contains the token “537.” If the value contains the specified token, then the second IF function returns an output of 1.0. If the value does not contain the specified token, then the second IF function returns an output of 0.0.
- (76) If the value in the column “foo.parameters.params.vndr” does not contain the specified token, then the first IF function proceeds to a third IF function. The third IF function determines whether the value of the input data in the column “fooElement.barHash” is equal to “77957379.0.” If the value is equal to the specified amount, then the third IF function returns an output of 1.0. If the value is not equal to the specified amount, then the third IF function returns an output of 0.0.
- (77) An example portion of program instructions corresponding to the decision node may be: if (type_BrowserSignal::bar.pv.nav[“fooKey”] defaults “default value”) contains “537” then 1.0 else 0.0)
- (78) In the above example, assume “type_BrowserSignal” corresponds to a set of input data, and “type_BrowserSignal::bar.pv.nav[“fooKey”]” references the variable corresponding to particular feature of the input data. The if-statement determines whether “type_BrowserSignal::bar.pv.nav[“fooKey”]” contains “537.” If the variable contains the specified token, then the if-statement returns a value of 1.0. If it does not contain the specified token, then the if-statement returns a value of 0.0. Additionally, “type_BrowserSignal::bar.pv.nav[“fooKey”]” defaults to a default value if the variable does not exist or the value is null.
- (79) In an embodiment, program instructions corresponding to a decision tree comprises a plurality of nested if-statements. Additionally, program instructions corresponding to a random forest may comprise a plurality of sets of program instructions, each set of program instructions corresponding to a particular decision tree. The output of the plurality of sets may be averaged to generate output of the random forest.
- (80) An example nested if-statement may be: if (lib Signal::foo.parameters.params.rndr contains “default” then (if (type_BrowserSignal::bar.pv.nav[“fooKey”] contains “537” then 1.0 else 0.0)) else (if (type_BrowserSignal::br.foo.barHash==77957379.0 then 1.0 else 0.0)))
- (81) In the above example, a first if-statement determines whether the variable “foo.parameters.params.rndr” contains a default value string token. If the variable contains the specified token, then the first if-statement proceeds to a second if-statement. As discussed above,

the second if-statement determines whether “bar.pv.nav[“fooKey”]” contains “537.” If the variable contains the specified token, then the second if-statement returns a value of 1.0. If the variable does not contain the specified token, then the second if-statement returns a value of 0.0.

(82) If the variable “foo.parameters.params.rndr” does not contain the specified token, then the first if-statement proceeds to a third if-statement. The third if-statement determines whether the variable “br.foo.barHash” is equal to “77957379.0.” If the variable is equal to the specified amount, then the third if-statement returns a value of 1.0. If the value is not equal to the specified amount, then the third IF function returns a value of 0.0.

(83) In an embodiment, the one or more program instructions or one or more queries may be executed on a system or computing device separate from system **300**. For example, system **300** may generate one or more program instructions and send the one or more program instructions to separate computing devices that execute the one or more program instructions. As another example, system **300** may generate one or more queries and send the one or more queries to a separate computing device.

(84) Additionally or alternatively, system **300** may receive input from a system or computing device separate from system **300**, execute the one or more program instructions or one or more queries based on the received input, and provide the output to the system or the computing device.

(85) For example, decision tree training module **310** may train random forest **320** to receive input comprising a request from a client computer to a server computer and determine whether the request is legitimate or illegitimate. As an example, a legitimate request may be from browser of a legitimate user, while an illegitimate request may be from a headless browser or a bot. A separate computing device, such as a server computer, may receive or intercept requests from client computing devices. The separate computing device may execute the one or more queries and/or one or more program instructions using data from the requests as input. Alternately, the computing device may send request data to system **300** and receive an indication from system **300** indicating whether the request was legitimate or illegitimate. Additionally or alternately, the computing device may store request data in a database at or accessible to system **300**, and cause system **300** to execute the one or more queries based on the stored request data.

4.2 Data Engine

(86) In the embodiment illustrated by FIG. 3, application programs interact with data stored in data store **340** via data engine **330**. Data engine **330** may be programmed or configured to perform a variety of functions needed for processing requests for data including, but not limited to: receiving queries from application program or decision tree training module **310**, executing queries, retrieving data from data store **340**, storing data to data store **340**, translating retrieved data into a requested format, sending query results to application program or decision tree training module **310**, and any other aspects of embodiments described herein.

(87) Additionally or alternatively, data engine **330** provides an API that may be accessed for example, by application program or decision tree training module **310**, to perform the functions related to processing data requests.

(88) In an embodiment, data engine **330** comprises a plurality of computing instances for processing requests. In response to receiving a request, data engine **330** selects one or more computing instances of the plurality to process the request. For example, in response to receiving a plurality of decision node training queries **312**, data engine **330** may execute the plurality of queries in parallel using a plurality of computing instances. Additionally or alternatively, each query of the plurality of decision node training queries **312** may be executed using a plurality of computing instances. By executing one or more decision node training queries in parallel, one or more decision trees may be trained more quickly compared to typical approaches that generate each layer of the decision tree and/or each decision node at a time, sequentially.

(89) Additionally or alternatively, data engine **330** may be programmed or configured to optimize query execution based on the requests received and/or the data stored at data store **340**. For

example, in response to receiving a plurality of decision node training queries **312**, data engine **330** may generate an execution plan that optimizes execution of the plurality of decision node training queries **312**.

4.3 Data Store

(90) In an embodiment, the data store **340** is a data storage subsystem consisting of programs and data that is stored on any suitable storage device such as one or more hard disk drives, memories, or any other electronic digital data recording device configured to store data. Although data storage system **340** is depicted as a single device in FIG. 3, data storage system **340** may span multiple devices located in one or more physical locations. In an embodiment, data storage system **340** is a distributed file system or a distributed data store comprising a plurality of data nodes. For example, data storage system **340** may be a Google Cloud Storage System, Hadoop Distributed File System (HDFS), MapR File System (MapR FS), HBase, or other suitable distributed data store.

(91) In an embodiment, data store **340** stores data in a column-oriented storage format. Example column-oriented data storage formats include Capacitor, Parquet, Kudu, and Object Relational Columnar (ORC). The particular data storage format may vary depending on the implementation.

(92) Generally, data store **340** stores data related to training a decision tree using a database system including, but not limited to, training input data **342** and result data **346**.

(93) Training input data **342** comprises one or more input sets for training a decision tree. For example, the training input data **342** may comprise a plurality of input sets corresponding to telemetry information for requests to a server computer from a plurality of client computers. Telemetry information may comprise data collected from a client computer that describes usages, states, behaviors, measurements, and/or other metrics, statistics, or values that may indicate what a client computer is, and/or is not, doing.

(94) The telemetry information may comprise one or more of: unprocessed data, processed data, or derived data. Unprocessed data comprises raw data collected from the client computer. Processed data may include, for example, converting data into string and/or numeric values, converting data into a standardized format, adding default values for missing attributes, etc. Derived data may be data derived from the collected data. For example, the collected data may include a first timestamp corresponding to a first keypress and a second timestamp corresponding to a second keypress, and the derived data may be a duration calculated based on the first and second timestamps.

(95) In an embodiment, the training input data **342** is stored in a columnar format. Each row may correspond to a particular input set of the training input, and each column may correspond to a particular feature of the training input. For example, assume input data comprises a vector, $\langle v1, v2, v3 \dots \rangle$. The vector may be stored in a particular row such that the value in column **1** is “v1,” the value in column **2** is “v2,” etc.

(96) One benefit of storing an input set in columnar format is that calculations may be performed for each feature in parallel. Referring to the above example, assume the input set comprises a large number of features, where each feature is stored in a respective column. The columns may be processed in parallel with one another. This results in significantly faster computation compared to processing each feature at a time for an input set.

(97) An additional benefit is that a database system is able to index the data stored in each column. For searching on a single field, only the column associated with the field needs to be loaded, and the indexing allows searching to be performed faster. Additionally, the database system may divide the columns into chunks. This allows the database system to identify which chunks are relevant for a computation and only operate on relevant chunks. The database system may also compress computations when performing calculations for each column. For example, if the cardinality of a column is low, that is if a column does not have many different values or many rows share the same values in a particular column, then a calculation may be performed just a single time for the rows.

(98) Result data **346** stores results generated by the plurality of decision tree training queries **312**. In an embodiment, result data **346** stores data indicating the variance associated with one or more

features of training input data **342**. Result data **346** may be used by decision tree training module **310** to determine which feature is associated with the highest variance. In an embodiment, result data **346** comprises a plurality of result sets, each result set corresponding to a particular decision node of a plurality of decision nodes being trained.

(99) In an embodiment, result data **346** comprises temporary tables. The temporary tables may be kept only for the duration of executing the decision node training queries **312**. In the illustrated embodiment, result data **346** is stored in data store **340**. Additionally or alternatively, the result data **346** may be stored in memory of one or more computing devices, such as memory of data engine **330** or one or more computing devices executing decision tree training module **310**.

(100) Additionally or alternatively, the data store **340** may be communicatively connected with other components, either directly or indirectly, such as one or more third party data suppliers, application program, or other computing devices. Data store **340** may be configured to receive or retrieve input data **342** and/or output data **344** from the one or more third party data suppliers, application program, and/or other computing devices.

5. Process Overview

(101) FIG. **4** illustrates example steps for training a decision tree using a database system.

(102) Although the steps in FIG. **4** are shown in an order, the steps of FIG. **4** may be performed in any order and are not limited to the order shown in FIG. **4**. Additionally, some steps may be optional, may be performed multiple times, and/or may be performed by different components.

(103) At step **400**, training data for training the decision tree is stored. The training data may comprise a plurality of sets of training input data. Each set of training input data may comprise a plurality of features of the input data. As an example, assume system **300** receives training data for training a decision tree that determines whether a request from a client computer is legitimate or illegitimate. Each set of training input data comprises telemetry information associated with a particular request. The training input data may be stored in data store **340**. Additionally, the training data may not include a plurality of training output data corresponding to the plurality of training input data.

(104) In an embodiment, each set of training input data is stored as a row of a database table. Each feature of the training input data may correspond to a particular row of the database table.

(105) Steps **402-408** are performed to generate a decision node of the decision tree. In some embodiments, the steps may be performed in parallel for some or all of the decision nodes. That is, some or all of the decision nodes may be generated at the same time.

(106) In an embodiment, a number of decision nodes for the decision tree is determined. The determination may be based on one or more factors such as a number of layers of the decision tree, a number of branches for each node of the decision tree, a structure of the decision tree, etc. For example, system **300** may be configured to generate decision trees that comprise 3 layers and 2 branches per node, for a total of up to 7 nodes. Based on the determined number of decision nodes, the below steps **402-408** are performed for each decision node. For example, if system **300** is configured to generate a decision tree with 7 nodes, decision tree training module **310** may repeat steps **402-408** 7 times.

(107) Additionally or alternatively, the decision nodes may be generated in serial, starting at a root node. Generating one or more subsequent decision nodes may be based on how previous decision nodes in the decision tree branches and/or the variance associated with features that have not been selected for previous decision nodes. As discussed above, after a decision node is generated, the input is divided or filtered based on the branches of the decision node and child nodes are generated based on the divided input data. Rather than generating a pre-determined number of decision nodes, decision tree training module **310** may iteratively repeat steps **402-408** until no features remain or the variance associated with all remaining features is below a particular threshold, e.g. 0.

(108) At step **402**, a set of training data is selected. The decision node is generated based on the

selected set of training data. In an embodiment, selecting a set of training data comprises selecting a subset of the plurality of sets of training input data and a corresponding subset of the output data. Additionally or alternatively, selecting a set of training data comprises selecting a subset of features from a plurality of features of the training input data.

(109) For the purpose of illustrating a clear example, assume training input data is stored in a database table. Each database row may correspond to a particular set of training input data, and each database column may correspond to a particular feature of the training input data. Selecting a set of training data may comprise selecting a plurality of rows from the database table and/or a plurality of columns, wherein the plurality of rows is fewer than all of the rows of the database table and the plurality of columns is fewer than all of the columns of the database table.

(110) At step **404**, the variance is calculated for each feature of a plurality of features of the selected training data using one or more queries. The plurality of features of the selected training data may be a subset of the features of the training input data.

(111) In an embodiment, the data type of each feature is determined. Based on the data type, a corresponding query is selected for calculating the variance of the feature. The corresponding query may be configured to parse data of the corresponding data type. Additionally, the corresponding query may select one or more feature values for which to calculate an associated variance. As discussed above, the method by which the one or more feature values is selected may vary based on the corresponding data type and the particular embodiment. The query calculates one or more variance values corresponding to the one or more selected feature values. The query may return a result indicating the highest variance and the feature value associated with the highest variance. Additionally or alternatively, the query may return a plurality of feature values and their associated variance. The results may be sorted by variance. Additionally, the results may be stored temporarily, such as in result data **346**.

(112) As an example, a particular feature may have string data values. The corresponding query may be configured to tokenize strings into a plurality of tokens, and for each token of the plurality, divide the input set based on the token, apply an aggregate label to the divided input sets, and calculate the variance associated with the token. The query may return a result indicating the token with the highest variance and the associated variance.

(113) At step **406**, the feature associated with the highest variance is identified. In an embodiment, identifying the feature associated with the highest variance comprises comparing a plurality of variance and feature (or feature value) pairs. Additionally or alternatively, identifying the feature associated with the highest variance comprises querying one or more results table. The query may retrieve, for each feature, the highest variance and the feature value associated with the highest variance, sort the results by variance, and return a result comprising the highest variance, feature, and/or the feature value associated with the highest variance.

(114) At step **408**, the decision node is generated based on the feature and feature value with the highest variance. The decision node branches on the feature value and the feature. For example, assume the feature associated with the highest variance is “screenDimensionX,” and the feature value of “screenDimensionX” associated with the highest variance is 1280. Decision tree training module **310** may generate a decision node that branches based on whether the value of the feature “screenDimensionX” is greater than 1280 or less than 1280.

(115) One branch of the decision node may correspond to a first result and the other branch may correspond to a second result. The first result and/or the second result may be a subsequent decision node. Additionally, the first result may be a first output value, and the second result may be a second output value. The first output value and the second output value may be based on the training output data.

(116) After all the decision nodes are generated, at step **410**, one or more queries corresponding to the decision tree are generated. The one or more queries may be written in any suitable query language, such as Structured Query Language (SQL) or other SQL-like languages. The particular

query language may vary depending on the implementation and the target computing system on which the one or more queries will be executed.

(117) Additionally or alternatively, one or more program instructions corresponding to the decision tree are generated. The one or more program instructions may be written in any suitable programming language. Additionally, the one or more program instructions may be compiled to generate one or more intermediate instructions that may be interpreted or executed by a computing device. The particular programming language may vary depending on the implementation and the target computing system on which the one or more program instructions will be executed.

(118) The one or more queries and/or the one or more program instructions, when executed, receive or retrieve a set of input data, apply the decision tree to the set of input data, and generate output indicating the result of applying the decision tree to the set of input data. The output may be, for example, a binary result such as I/O or true/false, or a categorical result, such as a particular classification from a set of possible classifications. In an embodiment, the output may be stored to a database table, such as in data store **340**, or to one or more data files that may be sent to or accessible by a computing device executing the one or more program instructions and/or the one or more queries. Example data files include JSON, comma separate values, SQL dumps, spreadsheets such as Excel files, and other file types. Additionally or alternatively, the output may be provided to another application, program, function, or module. Additionally or alternatively, the output may be displayed to a user of an application or program, such as a database client application.

(119) The one or more queries and/or the one or more program instructions may be provided to a computing device for utilizing the decision tree. Referring to the above example, an application may receive requests from a client computer to a server computer, execute the one or more queries and/or the one or more program instructions using request data as input, and receive output indicating whether the request is legitimate or illegitimate. The application may determine one or more actions to take based on whether the output indicates the request is legitimate or illegitimate.

(120) In an embodiment, generating one or more queries corresponding to the decision tree comprises generating, for each decision node of the decision tree, a query or a portion of a query corresponding to the decision node. The query or the portion of the query corresponding to the decision node may test a particular input data feature. The query, or portion of the query, when executed, may return a first value if the test evaluates to true, or a second value if the test evaluates to false. The first value and/or the second value may be a value returned by a subsequent decision node. That is, rather than returning an output of, for example 1.0 or 0.0, the return value may be a second query or a second portion of a query corresponding to another decision node. Additionally or alternatively, the query, or portion of the query, when executed, may test a plurality of conditions and return a different value based on which condition of the plurality of conditions evaluates to true.

(121) In an embodiment, generating one or more program instructions corresponding to the decision tree comprises generating, for each decision node of the decision tree, one or more respective program instructions corresponding to the decision node. The program instructions corresponding to the decision node may test a particular input data feature. The program instructions, when executed, may return a first value if the test evaluates to true, or a second value if the test evaluates to false. The first value and/or the second value may be a value returned by a subsequent decision node. The return value may comprise one or more second program instructions corresponding to another decision node.

(122) The steps **402-410** described above may be performed repeatedly to generate a plurality of decision trees for a random forest. In some embodiments, the steps may be performed in parallel for some or all of the decision nodes of the decision trees of the random forest. That is, some or all of the decision nodes may be generated at the same time, and some or all of the decision trees may be generated at the same time.

(123) In an embodiment, a number of decision trees for the random forest is determined. Each tree

may be generated using one or more techniques described above for generating a decision tree. In some embodiments, the steps may be performed in parallel for some or all of the decision nodes of the decision trees of the random forest. That is, some or all of the decision nodes may be generated at the same time, and some or all of the decision trees may be generated at the same time.

(124) In an embodiment, a number of decision trees is determined. Additionally, a number of decision nodes for each decision tree may be determined. The determination may be based on one or more factors such as a number of layers of each decision tree, a number of branches for each node of the decision tree, a structure of the decision tree, etc. The number of layers, number of branches, and/or the structure of each decision tree may be the same or different, depending on the implementation.

(125) As an example, assume system **300** is configured to generate 3 decision trees, where each decision tree comprises 3 layers and 2 branches per node, for a total of up to 21 decision nodes. Based on the determined number of decision nodes, the above steps **402-408** are performed for each decision node. Referring to the above example, if system **300** is configured to generate a random forest with 3 decision trees each with 7 nodes, decision tree training module **310** may repeat steps **402-408** 21 times.

(126) Additionally or alternatively, as discussed above, decision nodes of a decision tree may be generated in serial, where generating one or more subsequent decision nodes is based on how previous decision nodes in the decision tree branches and/or the variance associated with features that have not been selected for previous decision nodes. Although decision nodes of a decision tree are generated in serial, a plurality of decision trees may be generated in parallel.

(127) In an embodiment, after all the decision nodes are generated, one or more queries corresponding to the random forest are generated. Additionally or alternatively, one or more program instructions corresponding to the random forest are generated. The one or more queries and/or the one or more program instructions, when executed, receive or retrieve a set of input data, apply each decision tree of the random forest to the set of input data, and generate output indicating the result of applying the random forest to the set of input data.

(128) In an embodiment, generating the one or more queries and/or the one or more program instructions corresponding to the random forest comprises generating, for each decision tree of the random forest, one or more queries and/or one or more program instructions corresponding to the decision tree. A query and/or program instruction(s) corresponding to the random forest may include, or reference, the one or more queries and/or one or more program instructions corresponding to each decision tree. Additionally or alternatively, the query and/or program instruction(s) corresponding to the random forest may retrieve output generated by the one or more queries and/or one or more program instructions corresponding to each decision tree. The query and/or program instruction(s) corresponding to the random forest may process the outputs corresponding to the decision trees in the random forest and generate an output for the random forest. For example, a majority-vote of the output of the decision trees may be determined, or an average of the output of the decision trees may be calculated.

(129) To illustrate a clear example, assume the random forest comprises 3 decision trees, and each decision tree outputs a value of 1.0 or 0.0. An example set of program instructions corresponding to the random forest may be: $(\text{decisionTree1} + \text{decisionTree2} + \text{decisionTree3})/3$

where “decisionTree1” refers to the output generated by a first set of program instructions corresponding to a first decision tree, “decisionTree2” refers to the output generated by a second set of program instructions corresponding to a second decision tree, and “decisionTree3” corresponds to a third set of program instructions corresponding to a third decision tree. Each of “decisionTree1,” “decisionTree2,” and “decisionTree3” may be a call to the respective set of corresponding program instructions and/or to retrieve output generated by executing the respective set of corresponding program instructions. Alternately, each of “decisionTree1,” “decisionTree2,” and “decisionTree3” may be replaced in-line by the respective set of corresponding program

instructions.

6. Hardware Overview

(130) According to one embodiment, the techniques described herein are implemented by at least one computing device. The techniques may be implemented in whole or in part using a combination of at least one server computer and/or other computing devices that are coupled using a network, such as a packet data network. The computing devices may be hard-wired to perform the techniques, or may include digital electronic devices such as at least one application-specific integrated circuit (ASIC) or field programmable gate array (FPGA) that is persistently programmed to perform the techniques, or may include at least one general purpose hardware processor programmed to perform the techniques pursuant to program instructions in firmware, memory, other storage, or a combination. Such computing devices may also combine custom hard-wired logic, ASICs, or FPGAs with custom programming to accomplish the described techniques. The computing devices may be server computers, workstations, personal computers, portable computer systems, handheld devices, mobile computing devices, wearable devices, body mounted or implantable devices, smartphones, smart appliances, internetworking devices, autonomous or semi-autonomous devices such as robots or unmanned ground or aerial vehicles, any other electronic device that incorporates hard-wired and/or program logic to implement the described techniques, one or more virtual computing machines or instances in a data center, and/or a network of server computers and/or personal computers.

(131) FIG. 5 is a block diagram that illustrates an example computer system with which an embodiment may be implemented. In the example of FIG. 5, a computer system **500** and instructions for implementing the disclosed technologies in hardware, software, or a combination of hardware and software, are represented schematically, for example as boxes and circles, at the same level of detail that is commonly used by persons of ordinary skill in the art to which this disclosure pertains for communicating about computer architecture and computer systems implementations.

(132) Computer system **500** includes an input/output (I/O) subsystem **502** which may include a bus and/or other communication mechanism(s) for communicating information and/or instructions between the components of the computer system **500** over electronic signal paths. The I/O subsystem **502** may include an I/O controller, a memory controller and at least one I/O port. The electronic signal paths are represented schematically in the drawings, for example as lines, unidirectional arrows, or bidirectional arrows.

(133) At least one hardware processor **504** is coupled to I/O subsystem **502** for processing information and instructions. Hardware processor **504** may include, for example, a general-purpose microprocessor or microcontroller and/or a special-purpose microprocessor such as an embedded system or a graphics processing unit (GPU) or a digital signal processor or ARM processor. Processor **504** may comprise an integrated arithmetic logic unit (ALU) or may be coupled to a separate ALU.

(134) Computer system **500** includes one or more units of memory **506**, such as a main memory, which is coupled to I/O subsystem **502** for electronically digitally storing data and instructions to be executed by processor **504**. Memory **506** may include volatile memory such as various forms of random-access memory (RAM) or other dynamic storage device. Memory **506** also may be used for storing temporary variables or other intermediate information during execution of instructions to be executed by processor **504**. Such instructions, when stored in non-transitory computer-readable storage media accessible to processor **504**, can render computer system **500** into a special-purpose machine that is customized to perform the operations specified in the instructions.

(135) Computer system **500** further includes non-volatile memory such as read only memory (ROM) **508** or other static storage device coupled to I/O subsystem **502** for storing information and instructions for processor **504**. The ROM **508** may include various forms of programmable ROM (PROM) such as erasable PROM (EPROM) or electrically erasable PROM (EEPROM). A unit of persistent storage **510** may include various forms of non-volatile RAM (NVRAM), such as FLASH

memory, or solid-state storage, magnetic disk or optical disk such as CD-ROM or DVD-ROM, and may be coupled to I/O subsystem **502** for storing information and instructions. Storage **510** is an example of a non-transitory computer-readable medium that may be used to store instructions and data which when executed by the processor **504** cause performing computer-implemented methods to execute the techniques herein.

(136) The instructions in memory **506**, ROM **508** or storage **510** may comprise one or more sets of instructions that are organized as modules, methods, objects, functions, routines, or calls. The instructions may be organized as one or more computer programs, operating system services, or application programs including mobile apps. The instructions may comprise an operating system and/or system software; one or more libraries to support multimedia, programming or other functions; data protocol instructions or stacks to implement TCP/IP, HTTP or other communication protocols; file format processing instructions to parse or render files coded using HTML, XML, JPEG, MPEG or PNG; user interface instructions to render or interpret commands for a graphical user interface (GUI), command-line interface or text user interface; application software such as an office suite, internet access applications, design and manufacturing applications, graphics applications, audio applications, software engineering applications, educational applications, games or miscellaneous applications. The instructions may implement a web server, web application server or web client. The instructions may be organized as a presentation layer, application layer and data storage layer such as a relational database system using structured query language (SQL) or no SQL, an object store, a graph database, a flat file system or other data storage.

(137) Computer system **500** may be coupled via I/O subsystem **502** to at least one output device **512**. In one embodiment, output device **512** is a digital computer display. Examples of a display that may be used in various embodiments include a touch screen display or a light-emitting diode (LED) display or a liquid crystal display (LCD) or an e-paper display. Computer system **500** may include other type(s) of output devices **512**, alternatively or in addition to a display device. Examples of other output devices **512** include printers, ticket printers, plotters, projectors, sound cards or video cards, speakers, buzzers or piezoelectric devices or other audible devices, lamps or LED or LCD indicators, haptic devices, actuators or servos.

(138) At least one input device **514** is coupled to I/O subsystem **502** for communicating signals, data, command selections or gestures to processor **504**. Examples of input devices **514** include touch screens, microphones, still and video digital cameras, alphanumeric and other keys, keypads, keyboards, graphics tablets, image scanners, joysticks, clocks, switches, buttons, dials, slides, and/or various types of sensors such as force sensors, motion sensors, heat sensors, accelerometers, gyroscopes, and inertial measurement unit (IMU) sensors and/or various types of transceivers such as wireless, such as cellular or Wi-Fi, radio frequency (RF) or infrared (IR) transceivers and Global Positioning System (GPS) transceivers.

(139) Another type of input device is a control device **516**, which may perform cursor control or other automated control functions such as navigation in a graphical interface on a display screen, alternatively or in addition to input functions. Control device **516** may be a touchpad, a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **504** and for controlling cursor movement on display **512**. The input device may have at least two degrees of freedom in two axes, a first axis (e.g., x) and a second axis (e.g., y), that allows the device to specify positions in a plane. Another type of input device is a wired, wireless, or optical control device such as a joystick, wand, console, steering wheel, pedal, gearshift mechanism or other type of control device. An input device **514** may include a combination of multiple different input devices, such as a video camera and a depth sensor.

(140) In another embodiment, computer system **500** may comprise an internet of things (IoT) device in which one or more of the output devices **512**, input device **514**, and control device **516** are omitted. Or, in such an embodiment, the input device **514** may comprise one or more cameras, motion detectors, thermometers, microphones, seismic detectors, other sensors or detectors,

measurement devices or encoders and the output device **512** may comprise a special-purpose display such as a single-line LED or LCD display, one or more indicators, a display panel, a meter, a valve, a solenoid, an actuator or a servo.

(141) When computer system **500** is a mobile computing device, input device **514** may comprise a global positioning system (GPS) receiver coupled to a GPS module that is capable of triangulating to a plurality of GPS satellites, determining and generating geo-location or position data such as latitude-longitude values for a geophysical location of the computer system **500**. Output device **512** may include hardware, software, firmware and interfaces for generating position reporting packets, notifications, pulse or heartbeat signals, or other recurring data transmissions that specify a position of the computer system **500**, alone or in combination with other application-specific data, directed toward host **524** or server **530**.

(142) Computer system **500** may implement the techniques described herein using customized hard-wired logic, at least one ASIC or FPGA, firmware and/or program instructions or logic which when loaded and used or executed in combination with the computer system causes or programs the computer system to operate as a special-purpose machine. According to one embodiment, the techniques herein are performed by computer system **500** in response to processor **504** executing at least one sequence of at least one instruction contained in main memory **506**. Such instructions may be read into main memory **506** from another storage medium, such as storage **510**. Execution of the sequences of instructions contained in main memory **506** causes processor **504** to perform the process steps described herein. In alternative embodiments, hard-wired circuitry may be used in place of or in combination with software instructions.

(143) The term “storage media” as used herein refers to any non-transitory media that store data and/or instructions that cause a machine to operation in a specific fashion. Such storage media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage **510**. Volatile media includes dynamic memory, such as memory **506**. Common forms of storage media include, for example, a hard disk, solid state drive, flash drive, magnetic data storage medium, any optical or physical data storage medium, memory chip, or the like.

(144) Storage media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between storage media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise a bus of I/O subsystem **502**. Transmission media can also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(145) Various forms of media may be involved in carrying at least one sequence of at least one instruction to processor **504** for execution. For example, the instructions may initially be carried on a magnetic disk or solid-state drive of a remote computer. The remote computer can load the instructions into its dynamic memory and send the instructions over a communication link such as a fiber optic or coaxial cable or telephone line using a modem. A modem or router local to computer system **500** can receive the data on the communication link and convert the data to a format that can be read by computer system **500**. For instance, a receiver such as a radio frequency antenna or an infrared detector can receive the data carried in a wireless or optical signal and appropriate circuitry can provide the data to I/O subsystem **502** such as place the data on a bus. I/O subsystem **502** carries the data to memory **506**, from which processor **504** retrieves and executes the instructions. The instructions received by memory **506** may optionally be stored on storage **510** either before or after execution by processor **504**.

(146) Computer system **500** also includes a communication interface **518** coupled to bus **502**. Communication interface **518** provides a two-way data communication coupling to network link(s) **520** that are directly or indirectly connected to at least one communication networks, such as a network **522** or a public or private cloud on the Internet. For example, communication interface **518** may be an Ethernet networking interface, integrated-services digital network (ISDN) card,

cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of communications line, for example an Ethernet cable or a metal cable of any kind or a fiber-optic line or a telephone line. Network **522** broadly represents a local area network (LAN), wide-area network (WAN), campus network, internetwork or any combination thereof. Communication interface **518** may comprise a LAN card to provide a data communication connection to a compatible LAN, or a cellular radiotelephone interface that is wired to send or receive cellular data according to cellular radiotelephone wireless networking standards, or a satellite radio interface that is wired to send or receive digital data according to satellite wireless networking standards. In any such implementation, communication interface **518** sends and receives electrical, electromagnetic or optical signals over signal paths that carry digital data streams representing various types of information.

(147) Network link **520** typically provides electrical, electromagnetic, or optical data communication directly or through at least one network to other data devices, using, for example, satellite, cellular, Wi-Fi, or BLUETOOTH technology. For example, network link **520** may provide a connection through a network **522** to a host computer **524**.

(148) Furthermore, network link **520** may provide a connection through network **522** or to other computing devices via internetworking devices and/or computers that are operated by an Internet Service Provider (ISP) **526**. ISP **526** provides data communication services through a world-wide packet data communication network represented as internet **528**. A server computer **530** may be coupled to internet **528**. Server **530** broadly represents any computer, data center, virtual machine or virtual computing instance with or without a hypervisor, or computer executing a containerized program system such as DOCKER or KUBERNETES. Server **530** may represent an electronic digital service that is implemented using more than one computer or instance and that is accessed and used by transmitting web services requests, uniform resource locator (URL) strings with parameters in HTTP payloads, API calls, app services calls, or other service calls. Computer system **500** and server **530** may form elements of a distributed computing system that includes other computers, a processing cluster, server farm or other organization of computers that cooperate to perform tasks or execute applications or services. Server **530** may comprise one or more sets of instructions that are organized as modules, methods, objects, functions, routines, or calls. The instructions may be organized as one or more computer programs, operating system services, or application programs including mobile apps. The instructions may comprise an operating system and/or system software; one or more libraries to support multimedia, programming or other functions; data protocol instructions or stacks to implement TCP/IP, HTTP or other communication protocols; file format processing instructions to parse or render files coded using HTML, XML, JPEG, MPEG or PNG; user interface instructions to render or interpret commands for a graphical user interface (GUI), command-line interface or text user interface; application software such as an office suite, internet access applications, design and manufacturing applications, graphics applications, audio applications, software engineering applications, educational applications, games or miscellaneous applications. Server **530** may comprise a web application server that hosts a presentation layer, application layer and data storage layer such as a relational database system using structured query language (SQL) or no SQL, an object store, a graph database, a flat file system or other data storage.

(149) Computer system **500** can send messages and receive data and instructions, including program code, through the network(s), network link **520** and communication interface **518**. In the Internet example, a server **530** might transmit a requested code for an application program through Internet **528**, ISP **526**, local network **522** and communication interface **518**. The received code may be executed by processor **504** as it is received, and/or stored in storage **510**, or other non-volatile storage for later execution.

(150) The execution of instructions as described in this section may implement a process in the form of an instance of a computer program that is being executed, and consisting of program code

and its current activity. Depending on the operating system (OS), a process may be made up of multiple threads of execution that execute instructions concurrently. In this context, a computer program is a passive collection of instructions, while a process may be the actual execution of those instructions. Several processes may be associated with the same program; for example, opening up several instances of the same program often means more than one process is being executed. Multitasking may be implemented to allow multiple processes to share processor 504. While each processor 504 or core of the processor executes a single task at a time, computer system 500 may be programmed to implement multitasking to allow each processor to switch between tasks that are being executed without having to wait for each task to finish. In an embodiment, switches may be performed when tasks perform input/output operations, when a task indicates that it can be switched, or on hardware interrupts. Time-sharing may be implemented to allow fast response for interactive user applications by rapidly performing context switches to provide the appearance of concurrent execution of multiple processes simultaneously. In an embodiment, for security and reliability, an operating system may prevent direct communication between independent processes, providing strictly mediated and controlled inter-process communication functionality.

Claims

1. A computer-implemented method for training a decision tree using a database system, the decision tree comprising a plurality of decision nodes, the computer-implemented method implemented by a network traffic management system comprising one or more network traffic apparatuses, one or more computing devices, or server devices, the method comprising: storing in a database input data for training the decision tree, the input data comprising a plurality of feature values corresponding to a plurality of features, wherein the input data is stored in a columnar format, wherein rows of columnar format correspond to input sets of the input data in the database and columns correspond to at least one of the plurality of features, wherein the decision tree is trained using a set of training data comprising the input data, and wherein the set of training data does not comprise output data; and generating a particular node of the plurality of decision nodes by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data, wherein the calculated variance value is based on aggregate label output values of partitioned subsets of the input data that are generated based on identified patterns relating to legitimacy of requests to one of the server devices, wherein the partitioned subsets of the input data are based on one or more conditions, wherein the aggregate label output values provide status information on conformity of the partitioned subsets to the identified patterns relating to legitimacy of the requests; identifying a particular feature of the subset of the plurality of features associated with a highest variance value; and associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature associated with the highest variance value.
2. The method of claim 1 further comprising generating a random forest comprising a plurality of decision trees by generating, for each decision tree of the plurality of decision trees, a respective plurality of decision nodes.
3. The method of claim 2 further comprising generating a second query to the database system corresponding to the random forest, wherein the second query, when executed using the database system, applies each decision tree of the plurality of decision trees to a set of input data.
4. The method of claim 1 wherein the selecting the subset of the plurality of features comprises randomly selecting the subset of the plurality of features.
5. The method of claim 1 wherein calculating the variance value associated with the feature comprises: identifying the one or more conditions for partitioning the subset of the input data;

partitioning the subset of the input data based on the one or more conditions; applying an aggregate label to each partition of the subset of the input data; calculating a respective variance value for each condition of the one or more conditions based on applying the aggregate label to each partition of the subset of the input data; and selecting a highest respective variance value as the variance value associated with the feature.

6. A system for training a decision tree using a database system, the decision tree comprising a plurality of decision nodes, the system comprising: one or more processors; and non-transitory computer-readable media storing instructions which, when executed by the one or more processors, cause: storing in a database input data for training the decision tree, the input data comprising a plurality of feature values corresponding to a plurality of features, wherein the input data is stored in a columnar format, wherein rows of columnar format correspond to input sets of the input data in the database and columns correspond to at least one of the plurality of features, wherein the decision tree is trained using a set of training data comprising the input data, and wherein the set of training data does not comprise output data; and generating a particular node of the plurality of decision nodes by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data, wherein the calculated variance value is based on aggregate label output values of partitioned subsets of the input data that are generated based on identified patterns relating to legitimacy of requests to one of the server devices, wherein the partitioned subsets of the input data are based on one or more conditions, wherein the aggregate label output values provide status information on conformity of the partitioned subsets to the identified patterns relating to legitimacy of the requests; identifying a particular feature of the subset of the plurality of features associated with a highest variance value; and associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature associated with the highest variance value.

7. The system of claim 6 wherein the instructions, when executed by the one or more processors, further cause generating a random forest comprising a plurality of decision trees by generating, for each decision tree of the plurality of decision trees, a respective plurality of decision nodes.

8. The system of claim 7 wherein the instructions, when executed by the one or more processors, further cause generating a second query to the database system corresponding to the random forest, wherein the second query, when executed using the database system, applies each decision tree of the plurality of decision trees to a set of input data.

9. The system of claim 6 wherein the subset of the plurality of features comprises randomly selecting the subset of the plurality of features.

10. The system of claim 6 wherein calculating the variance value associated with the feature comprises: identifying the one or more conditions for partitioning the subset of the input data; partitioning the subset of the input data based on the one or more conditions; applying an aggregate label to each partition of the subset of the input data; calculating a respective variance value for each condition of the one or more conditions based on applying the aggregate label to each partition of the subset of the input data; and selecting a highest respective variance value as the variance value associated with the feature.

11. A non-transitory computer readable medium having stored thereon instructions for training a decision tree using a database system, the decision tree comprising a plurality of decision nodes, the non-transitory computer readable medium comprising executable code which when executed by one or more processors, causes the one or more processors to: store in a database input data for training the decision tree, the input data comprising a plurality of feature values corresponding to a plurality of features, wherein the input data is stored in a columnar format, wherein rows of columnar format correspond to input sets of the input data in the database and columns correspond to at least one of the plurality of features, wherein the decision tree is trained using a set of training

data comprising the input data, and wherein the set of training data does not comprise output data; and generate a particular node of the plurality of decision nodes by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data, wherein the calculated variance value is based on aggregate label output values of partitioned subsets of the input data that are generated based on identified patterns relating to legitimacy of requests to one of the server devices, wherein the partitioned subsets of the input data are based on one or more conditions, wherein the aggregate label output values provide status information on conformity of the partitioned subsets to the identified patterns relating to legitimacy of the requests; identifying a particular feature of the subset of the plurality of features associated with a highest variance value; and associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature associated with the highest variance value.

12. The medium as set forth in claim 11, wherein the executable code which when executed by the one or more processors, further causes the processors to: generate a random forest comprising a plurality of decision trees by generating, for each decision tree of the plurality of decision trees, a respective plurality of decision nodes.

13. The medium as set forth in claim 12, wherein the executable code which when executed by the one or more processors, further causes the processors to: generate a second query to the database system corresponding to the random forest, wherein the second query, when executed using the database system, applies each decision tree of the plurality of decision trees to a set of input data.

14. The medium as set forth in claim 11, wherein the selecting the subset of the plurality of features comprises randomly selecting the subset of the plurality of features.

15. The medium as set forth in claim 11, wherein calculating the variance value associated with the feature comprises: identifying the one or more conditions for partitioning the subset of the input data; partitioning the subset of the input data based on the one or more conditions; applying an aggregate label to each partition of the subset of the input data; calculating a respective variance value for each condition of the one or more conditions based on applying the aggregate label to each partition of the subset of the input data; and selecting a highest respective variance value as the variance value associated with the feature.

16. A network traffic manager apparatus for training a decision tree using a database system, comprising memory comprising programmed instructions stored in the memory and one or more processors configured to be capable of executing the programmed instructions stored in the memory to: store in a database input data for training a decision tree, the decision tree comprising a plurality of decision nodes, the input data comprising a plurality of feature values corresponding to a plurality of features, wherein the input data is stored in a columnar format, wherein rows of columnar format correspond to input sets of the input data in the database and columns correspond to at least one of the plurality of features, wherein the decision tree is trained using a set of training data comprising the input data, and wherein the set of training data does not comprise output data; and generate a particular node of the plurality of decision nodes by: selecting a subset of the plurality of features and a subset of the input data; using one or more queries to the database system, for each feature of the subset of the plurality of features, calculating a variance value associated with the feature based on the subset of the input data, wherein the calculated variance value is based on aggregate label output values of partitioned subsets of the input data that are generated based on identified patterns relating to legitimacy of requests to one of the server devices, wherein the partitioned subsets of the input data are based on one or more conditions, wherein the aggregate label output values provide status information on conformity of the partitioned subsets to the identified patterns relating to legitimacy of the requests; identifying a particular feature of the subset of the plurality of features associated with a highest variance value;

and associating the particular node with the particular feature, wherein the particular node causes the decision tree to branch based on the particular feature associated with the highest variance value.

17. The network traffic manager apparatus as set forth in claim 16, wherein the one or more processors are further configured to be capable of executing the programmed instructions stored in the memory to: generate a random forest comprising a plurality of decision trees by generating, for each decision tree of the plurality of decision trees, a respective plurality of decision nodes.

18. The network traffic manager apparatus as set forth in claim 17, wherein the one or more processors are further configured to be capable of executing the programmed instructions stored in the memory to: generate a second query to the database system corresponding to the random forest, wherein the second query, when executed using the database system, applies each decision tree of the plurality of decision trees to a set of input data.

19. The network traffic manager apparatus as set forth in claim 16, wherein the selecting the subset of the plurality of features comprises randomly selecting the subset of the plurality of features.

20. The network traffic manager apparatus as set forth in claim 16, wherein calculating the variance value associated with the feature comprises: identifying the one or more conditions for partitioning the subset of the input data; partitioning the subset of the input data based on the one or more conditions; applying an aggregate label to each partition of the subset of the input data; calculating a respective variance value for each condition of the one or more conditions based on applying the aggregate label to each partition of the subset of the input data; and selecting a highest respective variance value as the variance value associated with the feature.
