

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12407494
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Keidar; Ron et al.

Secure key delivery

Abstract

Approaches in accordance with various illustrative embodiments provide for the encryption of communications going into and out of a device, such as a chip or proprietary bus. The encryption can occur in a central Root-of-Trust (RoT), which can include agents for individual communication protocols to generate session keys used to encrypt communications for individual sessions, and the data can be sent to a crypto engine for the respective communication protocol. A key tunnel unit can be used to receive a wrapped session key over the public bus and then unwrap the key in hardware, then able to then transmit the unwrapped session key to the corresponding crypto engine without exposing the session key to software executing on the device outside the RoT. The receiving inline crypto engine can then use that session key to encrypt session data to be transmitted to a separate device or destination.

Inventors: Keidar; Ron (San Diego, CA), Hu; Xinxing (Shanghai, CN), Lee; Hye Su (Santa Clara, CA)

Applicant: Nvidia Corporation (Santa Clara, CA)

Family ID: 92637104

Assignee: Nvidia Corporation (Santa Clara, CA)

Appl. No.: 18/186442

Filed: March 20, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240305447 A1	Sep. 12, 2024

Foreign Application Priority Data

WO	PCT/CN2023/077970	Feb. 23, 2023
----	-------------------	---------------

Publication Classification

Int. Cl.: H04L9/08 (20060101); H04L9/14 (20060101)

U.S. Cl.:

CPC H04L9/0822 (20130101); H04L9/0825 (20130101); H04L9/14 (20130101);

Field of Classification Search

CPC: H04L (9/0822); H04L (9/0825); H04L (9/14); H04L (9/0897)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11968295	12/2023	Migault	N/A	H04L 9/0891
2008/0063209	12/2007	Jaquette	380/284	G06F 21/80
2021/0176632	12/2020	Liu	N/A	H04W 12/0471
2022/0141026	12/2021	Smith	713/181	H04L 9/0891

OTHER PUBLICATIONS

European Search Report issued in corresponding EP Application No. 24159411.8, dated Jul. 12, 2024. cited by applicant

Han et al., “Scalable and Secure Virtualization of HSM With ScaleTrust”, IEEE/ACM Transactions on Networking, vol. 31, No. 4, Aug. 2023, pp. 1595-1610, XP011947062, doi: 10.1109/TNET.2022.3220427. cited by applicant

Primary Examiner: Getachew; Abiy

Attorney, Agent or Firm: Hogan Lovells US LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) This application claims priority to PCT Application Serial No. PCT/CN2023/077970, filed Feb. 23, 2023, and entitled “Secure Key Delivery Over a Non-Secure Connection,” which is hereby incorporated herein in its entirety and for all purposes.

BACKGROUND

(2) As data security becomes increasingly important, there is movement toward encrypting most, if not all, of the communications going into, or out of, a computing device or system-on-chip (SoC). For various communications interfaces, a central Root-of-Trust (RoT) needs to send encryption keys to the crypto engines of those interfaces. The delivery of those keys raises various potential issues, as the crypto engines are typically located at the edge of the die (next to the pads or soldering points out of the silicon) and the distance from RoT is significant, making dedicated wires or a dedicated bus for key delivery too expensive in many instances. Sending the keys over a public bus can avoid the expense of a dedicated wire or bus, but increases the risk of attacks on a

bus optimized for performance and not security. Further, a micro-controller unit (MCU) managing the crypto engines is typically limited in its ability to perform a standard security protocol (e.g., TLS) with the RoT. The MCUs are also typically not trusted, so even if an MCU can handle secure delivery it should have no access to the value of those keys.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Various embodiments in accordance with the present disclosure will be described with reference to the drawings, in which:

(2) FIG. 1 illustrates an example system allowing for secure network communication between devices that can be used to implement aspects of various embodiments;

(3) FIG. 2 illustrates an example communication of session keys between devices, which can be used in accordance with various embodiments;

(4) FIGS. 3A and 3B illustrate an ethe format of a wrapped key and a key tunnel unit for processing such a wrapped key, in accordance with various embodiments;

(5) FIG. 4 illustrates an example flow of wrapping and session keys that can be produced in accordance with various embodiments;

(6) FIG. 5 illustrates components of an example system for transmitting wrapped session keys over a public interface, in accordance with at least one embodiment;

(7) FIG. 6 illustrates an example process for transmitting wrapped session keys over an interface that can be performed in accordance with at least one embodiment;

(8) FIG. 7 illustrates an example data center system, according to at least one embodiment;

(9) FIG. 8 is a block diagram illustrating a computer system, according to at least one embodiment;

(10) FIG. 9 is a block diagram illustrating a computer system, according to at least one embodiment;

(11) FIG. 10 illustrates a computer system, according to at least one embodiment;

(12) FIG. 11 illustrates a computer system, according to at least one embodiment;

(13) FIG. 12A illustrates a computer system, according to at least one embodiment;

(14) FIG. 12B illustrates a computer system, according to at least one embodiment;

(15) FIG. 13 illustrates exemplary integrated circuits and associated graphics processors, according to at least one embodiment;

(16) FIGS. 14A-14B illustrate exemplary integrated circuits and associated graphics processors, according to at least one embodiment;

(17) FIGS. 15A-15B illustrate additional exemplary graphics processor logic according to at least one embodiment;

(18) FIG. 16 illustrates a computer system, according to at least one embodiment;

(19) FIG. 17A illustrates a parallel processor, according to at least one embodiment;

(20) FIG. 17B illustrates a partition unit, according to at least one embodiment;

(21) FIG. 18 illustrates a multi-graphics processing unit (GPU) system, according to at least one embodiment;

(22) FIG. 19 illustrates a graphics processor, according to at least one embodiment; and

(23) FIG. 20 illustrates at least portions of a graphics processor, according to one or more embodiments.

DETAILED DESCRIPTION

(24) In the following description, various embodiments will be described. For purposes of explanation, specific configurations and details are set forth in order to provide a thorough understanding of the embodiments. However, it will also be apparent to one skilled in the art that the embodiments may be practiced without the specific details. Furthermore, well-known features

may be omitted or simplified in order not to obscure the embodiment being described.

(25) Approaches in accordance with various illustrative embodiments provide for the encryption of communications going into and out of a device, such as a chip or a proprietary bus. The encryption can occur in a central Root-of-Trust (RoT), such as may correspond to at least one secure processor, which can include or host one or more agents for individual communication protocols to generate session keys used to encrypt communications for individual sessions, and the data can be sent to a crypto engine for the respective communication protocol on the device. Because there can be a significant distance between the RoT and the crypto engine, it can be beneficial to use an existing public bus to send the data rather than using a dedicated bus or wire for key delivery. In order to minimize the likelihood of exposure of the key when sent over a public bus, each session key can be wrapped (e.g., encrypted) within the RoT using a wrapping key. A hardware block (as may correspond to or provide a “key tunnel unit” (KTU)) can be used to receive a wrapped session key over the public bus and then unwrap the key in hardware, being able to then transmit the unwrapped session key to the corresponding crypto-engine without exposing the session key to software executing on the device outside the RoT. The receiving inline crypto engine (or other inline encryption hardware) can then use that session key to encrypt session data to be transmitted to a separate device or destination.

(26) In order to establish a shared secret, an initial wrapping key can be sent over a public bus to the KTU, followed by a first wrapping key that is wrapped by the initial wrapping key (which may have been exposed due to transmission over the public bus). A second wrapping key can be sent that is wrapped using the first wrapping key, and the second wrapping key (once unwrapped) can become the current wrapping key to be used to wrap session keys. This process can continue for a series of wrapping keys sent before any session keys are sent, in order to reduce a likelihood that someone intercepting the initial wrapping key sent over the public bus will be able to intercept and unwrap any or all of the sequence of wrapping keys. Once one or more wrapping keys has been sent, a current (or most recently sent or received) wrapping key can be used to wrap or encrypt one or more session keys sent subsequent to the current wrapping key, and those encrypted session keys can then be transmitted over a public system bus. An IP block, or key tunnel unit, can then unwrap the session key(s) in hardware for use in encrypting data to be transmitted to another device. The wrapping keys can be rotated over time as well, such as to send a new wrapping key after a period of time or number of session keys.

(27) Variations of this and other such functionality can be used as well within the scope of the various embodiments as would be apparent to one of ordinary skill in the art in light of the teachings and suggestions contained herein.

(28) FIG. 1 illustrates a system **100** allowing for secure network communication between devices that can be used to implement aspects of various embodiments. In this example, a first computing device **102** is to securely exchange or communicate data with a second computing device **130** over at least one network. The computing devices can be similar or different devices, and may include any appropriate devices capable of performing operations such as encrypting, decrypting, transmitting, receiving, and/or processing data. Such devices may include, for example, servers, desktop computers, or compute instances, among other such options. These devices can include processors (such as one or more central processing units (CPUs) **108**, **138** or graphics processing units (GPUs) and memory that can connect to a central system bus **104**, **132** using a memory interface **106**, **140**, and can communicate with the network **120** using respective network interfaces **110**, **134**. The network(s) **120** may include any appropriate wired and/or wireless network, as may include the Internet, a cellular network, an ethernet, a local area network (LAN), or a peer-to-peer network, among other such options.

(29) These devices **102**, **130** can establish a secure communication path between them even if the path is across a public or otherwise unsecure network **120**. In at least one embodiment, this can include performing a control, secret, or key exchange. This exchange may occur in accordance with

a key management protocol supported by each device, which may be specified by the device **102** attempting to establish the secure communication path. The key management protocol can be managed by one or more trusted components that operate in a trusted zone or portion of the device typically referred to as a root of trust (RoT), which can correspond to a secure processor or set of secure components on each device. This can include, for example, inline secure hardware such as an inline crypto engine **112**, **136** on each device. These trusted components can exchange protocol information during a key establishment phase. After authentication of both devices, these devices can exchange data and configuration information, and can agree on a session key that is unique for a specific communication session. One of the devices can go to the data plane, which can encrypt and decrypt data, provide the session key to be used for the current session. In some embodiments, the data plane encryption component(s) will not be a separate component or have high throughput, although for low level encryption the data plane might be separate or capable of high throughput that might require its own space and hardware implementation. In at least one embodiment, an inline crypto engine managing the data plane will need to be very close to the location of the RoT, as may be embedded near the pins. The inline crypto engines **112**, **136** on the devices **102**, **130** can perform operations such as a protocol exchange or key exchange to establish keys within this root of trust.

(30) As mentioned, the distance to the inline crypto engine from an RoT can be important for at least some devices. If keys are to be delivered to something nearby, or over a short distance, then one or more additional wires could be used to transmit the data. For large, modern chips, however, the distance over which the keys may need to be transmitted can be significant, and might even involve separate dies in a single package. It is still beneficial to be able to protect the keys as they are transmitted from this root of trust to the inline encryption engine (such as an ICE-based key tunnel unit (KTU)), regardless of the distance the keys are to travel.

(31) One approach that can be used to secure data transmitted over a communication channel is to use a shared secret. FIG. 2 illustrates an example implementation **200** that can be used to securely manage such shared secrets in accordance with at least one embodiment. A public interface **214** between the two devices can support data plane encryption, where data to be transmitted is encrypted on a first device **202** then transmitted over the public interface **214** to a second device **220**, which has the shared secret and can then decrypt the data. In at least one embodiment, standard data plane protocols can be used to protect the data. The same, or separate, communication channel **212** can be used to exchange information relating to the key management protocols, such as to specify on or agree to the (e.g., standard) key management protocols to be used to protect the session keys to be used for one or more sessions.

(32) In this example, a key management protocol component **204** on the initiating computing device **202** is to provide for key delivery within the computing device **202**, such as within a system on chip (SoC). This communication can involve communicating a session key to an inline crypto engine **206** that may be a significant distance from the key management protocol component **204** on the device **202**. In this example, a (virtual) key tunnel **210** can be used to protect the key delivery inside the SoC, where those keys are to be used to encrypt or decrypt data within the device **202**. The key tunnels **210**, **224** can support a proprietary or non-public implementation. A key tunnel **210** can be responsible for moving session keys from the root of trust, such as a trusted agent that supports the key management protocol **204**, to a remote engine, such as an inline crypto engine **206**.

(33) As mentioned the distance between the RoT and the ICE may be significant, so it may be undesirable to add or use one or more dedicated wires for the key tunnel. Using a publicly accessible bus does not provide adequate trust when delivering keys, particularly keys in plaintext. Further, the ICE **206** may use a local controller or microcontroller that is not trusted and, as such, should not have access to the session keys or other shared secrets. An approach in accordance with at least one embodiment can prevent the microcontroller from having access to the keys in

unencrypted form, and can instead ensure that the microcontroller only receives encrypted keys over the key tunnel and can push those encrypted keys into the ICE **206**, or other encryption hardware, without having access to the unencrypted session keys. The root of trust will typically have the ability to encrypt keys, and the ICE can receive the encrypted keys through the key tunnel unit **364** (in hardware), and can unwrap and decrypt the keys in such a way that the key in plaintext is only visible to the ICE engine in the hardware, without the keys being exposed to the software. In at least one embodiment, the plaintext key will be sent very early during an early execution process (e.g., boot or production), such that its exposure will be limited because there should not be any non-boot-related applications or processes running on the device at that point. In an embodiment using an RTL key, for example, there may be a pre-shared secret stored such that there is no need to send an initial plaintext key.

(34) In at least one embodiment, the hardware can inspect (e.g., passively check the correctness of a software input) the behavior of the software, and if the software attempts to perform an unauthorized operation such as to skip a key, insert a key twice, or change a key, for example, the hardware will detect the attacks and ensure that it will not accept those modified keys through modification attacks. The root of trust can therefore determine that it can trust the key tunnel hardware because the local microcontroller is not permitted to access or modify the keys or other secret information, and that the RoT will also be able to determine any time keys are modified, one or more counters **358** are incremented, or other such actions occur. The counters can be used in at least one embodiment to ensure that key identifiers are monotonically increasing, where any gaps or reuse of key identifiers can be indicative of a wrapped key being reused or skipped. A key tunnel **224** can perform similar operations on the second device **220**, providing a secure path between an agent (supporting the key management protocol **222**) of the root of trust and an ICE **226** on the device. As mentioned, the keys can be used to encrypt data to be transmitted over the public interface **214** using data encryption endpoints **208**, **228** on the devices that are parties to the communication.

(35) In at least one embodiment, the root of trust includes an IP block that can unwrap keys in hardware and send them to the hardware of the crypto engine without exposing the keys to the software executing on the device. Session keys can be wrapped by the root of trust using wrapping keys, and this wrapping can be performed in a trusted environment. Because the session keys will be wrapped, there is no need for a dedicated bus between the root of trust and the encryption hardware. By wrapping the keys at the RoT and enabling the keys to be unwrapped at the key tunnel **210**, the wrapped keys can be delivered on public busses without risk of exposure. The respective microcontroller unit (MCU) does not need to perform any complex operations, as these operations can be offloaded to the key tunnel. The key tunnel can isolate the keys and is not dependent on the security of the MCU, such that the encryption approach can be supported even for an untrusted MCU. A key tunnel can offer other benefits as well, as there may be different ways to initialize shared secrets, and a key tunnel can provide built-in protection and detection against a variety of attacks. A key tunnel can also provide a mechanism to deliver secret keys or entropy to the MCU, as in some cases the MCU may need one or more secrets or entropy (e.g., random values) in addition to those sent to the crypto engine. A key tunnel can also provide a way to deliver authenticated commands to different hardware blocks, as well as a way to rotate wrapping keys and handover ownership, such as for RoT management. A key tunnel can also occupy only a small hardware area, which can be important for many devices where real estate is at a premium.

(36) FIG. **3A** illustrates an example of a wrapped key payload **300** that can be generated in accordance with at least one embodiment. In this example, it can be seen that a portion of the payload is not encrypted but is integrity protected, as may include the manifest, controls, and metadata. An encrypted portion of the payload may include the wrapped key and any necessary padding, as well as any MAC or initialization vector (IV) data to be transmitted.

(37) FIG. **3B** illustrates an example key tunnel unit implementation **350** that can be utilized in

accordance with various embodiments. A key tunnel unit (KTU) **364**, which may include or provide an inline crypto engine, can support various secure operations, as may include inserting a plaintext wrapping key via serial interface or a register, as well as inserting a wrapped wrapping key with an RTL key or key from fuses. Such a block can also support the rotation of a wrapping key, as well as the insertion of a wrapped session key. The key tunnel unit can correspond to a piece of hardware at a data encryption endpoint, proximate the ICE, which in a least one embodiment can receive wrapped keys, unwrap or decrypt those keys, and then push those keys (such as unwrapped session keys **374**) to the ICE or other encryption hardware. This unwrapping and providing of the session key in hardware can occur without exposing the unwrapped session key to software **352** on the respective computing device.

(38) In at least one embodiment, a primary function of such a key tunnel unit (KTU) is to deliver session keys for use on corresponding sessions. There may be many such session keys, provided by large agents that manage keys for many different sessions. There might be a separate session for each destination, and for each destination there may be multiple virtual destinations or applications. A key unwrap engine **364** of the KTU can be responsible for unwrapping many wrapped keys **354** and providing them to the appropriate crypto engine associated with the appropriate session. As illustrated, the key tunnel unit **364** can receive many types of keys, as may include plaintext keys **356** from the software **352**, as well as other keys such as RTL keys **362** and the like. The block **364** can maintain a status **366**, **368** of the current wrapping key(s) and session key(s), and can use the appropriate current wrapping key **360** to unwrap a wrapped key **354** using an appropriate unwrap block **372** or operation to extract the session key. By receiving only wrapped session keys **354**, the ability for an unintended party to obtain the session key from the transmission will be minimized. In at least one embodiment, a given KTU will only store the current wrapping key, unless there are multiple current wrapping keys to be stored to use to wrap different sets of session keys. When a wrapped key is received, that key passes through the unwrap engine **372**, which is the engine that performs the unwrapping and verification of the tag or signature. After a key is unwrapped and verified, the determination is made as to whether that key is a session key **374** or a new wrapping key **360**. If it is a session key **374**, that key is transmitted to/through the ICE to the hardware that is integrated with the key tunnel. If the key is a new wrapping key **360**, the key will go back and replace the current wrapping key.

(39) In order to further minimize the likelihood of a session key being intercepted or otherwise obtained by an unintended party or recipient, there can be a number of additional steps taken to provide additional security. This may include rotation of the wrapping keys and session keys, for example, but can also include the transmission and usage of a sequence of keys for the purpose of making it difficult for an unintended party to receive and decrypt all relevant keys. FIG. 4 illustrates example flows **400**, **450** of keys that can be provided in accordance with various embodiments. In this example, there can be an initial key **402** (the fundamental shared secret agreed upon per the protocol) sent that is unwrapped. This can be a plaintext key, for example, or other keys such as a FUSE key, Elliptic-curve Diffie-Hellman (ECDHE) key, register transfer level (RTL) key, or serial interface key, among other such options. This key can be sent over a public bus or interface to the ICE that can be stored as a current key associated with a session.

(40) In this example, a first wrapping key can be sent (with an identifier such as SW_ID=1) that can be used to wrap one or more session keys **406** for one or more sessions. In order to reduce the likelihood of an unintended party obtaining the wrapping key **404**, the wrapping key can be wrapped with the initial key **402**, such as a received plaintext key. This wrapping would require an unintended party to intercept both the plaintext key and the wrapping key to be able to unwrap the wrapping key, then use that wrapping key to unwrap the session key. Further protection can be provided by, although now illustrated in FIG. 4, sending a series of wrapping keys that are each wrapped by the prior wrapping key. In this way, unless an unintended party is able to intercept or obtain all of these keys, the unintended party will not be able to unwrap the session key(s) that are

wrapped with the final wrapping key in that sequence.

(41) In the example illustrated in FIG. 4, once a wrapping key **404** is sent that is to be used to wrap session keys, and that wrapping key is unwrapped using any other prior wrapping key, that wrapping key **404** can be used to wrap one or more session keys **406**. These may be session keys for sessions associated with the same or different recipients in various embodiments. The KTU (or other hardware) can unwrap these session keys using the current wrapping key, and can provide the unwrapped session keys to the ICE.

(42) In at least one embodiment, it will be desirable to rotate or otherwise change the wrapping key used over time, such that if an unintended party obtains a wrapping key they will likely only be able to unwrap a limited number of session keys without also obtaining the subsequent new wrapping key, which will also be wrapped. The rotation or change can come at any appropriate time, such as after a specified length of time or a number of wrapped session keys, or upon manual instruction, among other such options. In this case, a new wrapping key **408** can be sent through the key tunnel that is wrapped using the prior wrapping key **404**. Once unwrapped, the new wrapping key **408** can be used to wrap one or more subsequent session keys **410**. In at least one embodiment, a sequence of wrapped wrapping keys may be sent for each rotation as well. The process can continue as illustrated, with new wrapping keys sent as part of a key rotation process to encrypt one or more session keys. As mentioned, there may be many different wrapping keys used at any given time, such that the wrapping key used for session keys of one user or party cannot be used to unwrap session keys for another user or party. As illustrated in the example key delivery sequence **450**, a sequence of wrapping keys can be received by the KTU, where each wrapping key can wrap session keys until a new wrapping key is received that is to be used to wrap one or more subsequent session keys. Each individual session key can be unique in this example, and session keys may either be required to each be unique and may be rotated after an amount of time or other such factor. In at least one embodiment, each key can have a flag (or other indicator) set to indicate whether that key is to be used as a session key or a wrapping key. A session key can also have associated metadata that indicates, for example, the appropriate session to which that key belongs, whether it replaces a prior session key for that session, or other such information. In at least some embodiments, this metadata will not be encrypted but can still be authenticated so that, for example, intermediate controllers can view but not change the metadata. In at least one embodiment, the software can control how often new wrapping keys are generated and used, using an exposed hardware option.

(43) FIG. 5 illustrates an example system **500** that can be used to store session keys in at least one embodiment. As illustrated, the root of trust hardware **502** (e.g., a secure processor on chip) may operate or host several agents **502**, **504**, **506**, **508** as discussed previously, which may be tasked with managing keys for various sessions, in cooperation with an KTU manager **510** on the software side, as well as other components (e.g., a crypto accelerator **512** and Key Mover component **514**) on the hardware side. The keys (wrapped or otherwise), key metadata, and other relevant data can be transmitted over at least one public interface **520** or bus (or at least partially unsecure interface or bus) to be received to the respective recipients, which in this case can include a PCI Express-Integrity and Data Encryption (PCIe-IDE) component **530**, a storage crypto engine **540**, and a MacSEC inline engine **550**, among other such options (e.g., other crypto hardware). In this example the session keys (stored in plaintext form in the key tables) are stored in respective symmetric key tables, using the keys unwrapped by the components (e.g., **364** in FIG. 3B) and the associated authenticated key metadata. In this example, each KTU block or component would store the current respective wrapping key(s) to use to unwrap corresponding session keys. In this example, the agents create manage the session keys while the KTU manager **510** manages the wrapping keys within the root of trust **502**. The KTU manager can also be responsible for rotating the wrapping keys and performing the wrapping in at least one embodiment, although in other embodiments the wrapping keys may be provided to the individual agents for the agents to wrap

the session keys. Agents for other or alternative protocols can be used as well within the scope of various embodiments. For an example protocol option **530**, a small microcontroller (RISCV) may perform some of the management, such as to read the metadata and use it for routing the key to the correct location in the key table. In the second case, the index is read directly by the hardware as a pure index, which can be used to control the routing of the key. For the MacSEC engine **550**, there is some hardware logic that analyzes the metadata and not the pure index, such as to determine a session ID or key type flag. The logic can convert or translate this metadata using a local indexing mechanism. Other approaches to processing wrapped keys received over a public interface (e.g., an on-chip system bus) can be used as well within the scope of various embodiments.

(44) In at least one embodiment there may be a handover mechanism and multiple roots of trust. A first root of trust might start the initialization to generate the first wrapping key or perform the first key rotation. The actual session keys might, however, be handled by a different root of trust. The handover mechanism can be a secure hardware mechanism to handover ownership of the key tunnels between roots of trust by, for example, handing over the current wrapping key. The root of trust receiving the handover can then function as the new or current root of trust that can send session keys. Upon taking ownership, that root of trust might first rotate the wrapping key to a new wrapping key, such that the prior root of trust does not know the shared secret and cannot do anything until it again obtains ownership. Such functionality can be important, as there may be different roots of trust for the system at different times, such as during boot versus at runtime. Further, such an approach can be important for operations involving confidential compute, for example, where only the current root of trust is to be able to access certain information. The current wrapping key can thus be passed between secure processors, which can take over as the current root of trust. Such an approach can be used to encrypt any or all communications going in and out of a chip, including DRAM, PCIe, Video, ethernet, and proprietary busses like NVLINK and Chip2Chip, among other such options.

(45) FIG. **6** illustrates an example process **600** for securely transmitting session keys over a public interface that can be performed in accordance with various embodiments. It should be understood that for this and other processes presented herein that there may be additional, fewer, or alternative steps performed or similar or alternative orders, or at least partially in parallel, within the scope of the various embodiments unless otherwise specifically stated. Further, although this example is described with respect to session keys, various other types of data may be secured advantageously for similar or other purposes as well within the scope of the various embodiments. In this example, a wrapping key is provided **602** to a key tunnel unit (KTU) component of cryptographic hardware on a computing device. As discussed previously herein, this may include one of a sequence of wrapping keys that were also each encrypted with a prior wrapping key or other shared secret, among other such options. A session key can then be generated **604** using a software agent executing on a secure processor (or otherwise within a root of trust) on the computing device. There may be several agents each managing session keys for different sessions in at least one embodiment. The session key can then be encrypted **606**, or “wrapped,” on the secure processor using the current wrapping key. The wrapped session key can then be transmitted **608** over a public bus (or other at least partially unsecure interface) of the computing device.

(46) In this example, the wrapped session key can be received to the key tunnel unit of the crypto hardware. The wrapped session key can be decrypted **610** or “unwrapped” using the current wrapping key stored by the key tunnel unit. The unwrapped session key can then be delivered **612** to an inline crypto engine, for example, on the computing device to allow the crypto engine to encrypt corresponding session data using the unwrapped (e.g., plaintext) session key. A determination can be made **614** as to whether the wrapping key should be rotated or otherwise changed to a new wrapping key. If not, the process can continue with the next session key being wrapped using the current wrapping key. If the wrapping key is to be rotated, then a new wrapping key can be provided **616** from a wrapping key manager to the key tunnel unit to be used as the

current wrapping key. The key tunnel unit can analyze a flag or value for the key to determine that the key is to be treated as a new wrapping key and not another session key. The process can then continue using the “new” current wrapping key. In other embodiments, a new wrapping key may not be generated until after a new session key is generated or received, and that new session key is to be transmitted using a new wrapping key, where by the new wrapping key will be transmitted first, then used to wrap the new session key to be transmitted. Various other orderings or approaches can be used as well within the scope of the various embodiments.

(47) In at least one embodiment, data to be transferred using these data connections between peer devices can be received from, or transmitted to, any number of a variety of computing systems, devices, and services, and this functionality can be implemented in various types of systems or networked environments. For example, transmitting client devices can include any appropriate computing devices, as may include a desktop computer, notebook computer, set-top box, streaming device, gaming console, smartphone, tablet computer, VR headset, AR goggles, wearable computer, or a smart television. Each client device can submit a request across at least one wired or wireless network, as may include the Internet, an Ethernet, a local area network (LAN), or a cellular network, among other such options, which may include a variety of interconnected network devices. In one example, these requests can be submitted to an address associated with a cloud provider, who may operate or control one or more electronic resources in a cloud provider environment, such as may include a data center or server farm. In at least one embodiment, the request may be received or processed by at least one edge server, that sits on a network edge and is outside at least one security layer associated with the cloud provider environment.

(48) In at least one embodiment, such a system can be used for performing graphical rendering operations. In other embodiments, such a system can be used for other purposes, such as for providing image or video content to test or validate autonomous machine applications, or for performing deep learning operations. In at least one embodiment, such a system can be implemented using an edge device, or may incorporate one or more Virtual Machines (VMs). In at least one embodiment, such a system can be implemented at least partially in a data center or at least partially using cloud computing resources. Such a system can also be used to generate or present virtual reality (VR), augmented reality (AR), or enhanced reality (ER) content. In other embodiments, such a system can be used to perform hardware testing using simulation, to generate synthetic data, or as part of a collaborative content creation platform for 3D assets, among other such options.

(49) Data Center

(50) FIG. 7 illustrates an example data center **700**, in which at least one embodiment may be used. In at least one embodiment, data center **700** includes a data center infrastructure layer **710**, a framework layer **720**, a software layer **730** and an application layer **740**.

(51) In at least one embodiment, as shown in FIG. 7, data center infrastructure layer **710** may include a resource orchestrator **712**, grouped computing resources **714**, and node computing resources (“node C.R.s”) **716(1)-716(N)**, where “N” represents a positive integer (which may be a different integer “N” than used in other figures). In at least one embodiment, node C.R.s **716(1)-716(N)** may include, but are not limited to, any number of central processing units (“CPUs”) or other processors (including accelerators, field programmable gate arrays (FPGAs), graphics processors, etc.), memory storage devices **718(1)-718(N)** (e.g., dynamic read-only memory, solid state storage or disk drives), network input/output (“NW I/O”) devices, network switches, virtual machines (“VMs”), power modules, and cooling modules, etc. In at least one embodiment, one or more node C.R.s from among node C.R.s **716(1)-716(N)** may be a server having one or more of above-mentioned computing resources.

(52) In at least one embodiment, grouped computing resources **714** may include separate groupings of node C.R.s housed within one or more racks (not shown), or many racks housed in data centers at various geographical locations (also not shown). In at least one embodiment, separate groupings

of node C.R.s within grouped computing resources **714** may include grouped compute, network, memory or storage resources that may be configured or allocated to support one or more workloads. In at least one embodiment, several node C.R.s including CPUs or processors may grouped within one or more racks to provide compute resources to support one or more workloads. In at least one embodiment, one or more racks may also include any number of power modules, cooling modules, and network switches, in any combination.

(53) In at least one embodiment, resource orchestrator **712** may configure or otherwise control one or more node C.R.s **716(1)-716(N)** and/or grouped computing resources **714**. In at least one embodiment, resource orchestrator **712** may include a software design infrastructure (“SDI”) management entity for data center **700**. In at least one embodiment, resource orchestrator **712** may include hardware, software or some combination thereof.

(54) In at least one embodiment, as shown in FIG. 7, framework layer **720** includes a job scheduler **722**, a configuration manager **724**, a resource manager **726** and a distributed file system **728**. In at least one embodiment, framework layer **720** may include a framework to support software **732** of software layer **730** and/or one or more application(s) **742** of application layer **740**. In at least one embodiment, software **732** or application(s) **742** may respectively include web-based service software or applications, such as those provided by Amazon Web Services, Google Cloud and Microsoft Azure. In at least one embodiment, framework layer **720** may be, but is not limited to, a type of free and open-source software web application framework such as Apache Spark™ (hereinafter “Spark”) that may utilize distributed file system **728** for large-scale data processing (e.g., “big data”). In at least one embodiment, job scheduler **722** may include a Spark driver to facilitate scheduling of workloads supported by various layers of data center **700**. In at least one embodiment, configuration manager **724** may be capable of configuring different layers such as software layer **730** and framework layer **720** including Spark and distributed file system **728** for supporting large-scale data processing. In at least one embodiment, resource manager **726** may be capable of managing clustered or grouped computing resources mapped to or allocated for support of distributed file system **728** and job scheduler **722**. In at least one embodiment, clustered or grouped computing resources may include grouped computing resources **714** at data center infrastructure layer **710**. In at least one embodiment, resource manager **726** may coordinate with resource orchestrator **712** to manage these mapped or allocated computing resources.

(55) In at least one embodiment, software **732** included in software layer **730** may include software used by at least portions of node C.R.s **716(1)-716(N)**, grouped computing resources **714**, and/or distributed file system **728** of framework layer **720**. In at least one embodiment, one or more types of software may include, but are not limited to, Internet web page search software, e-mail virus scan software, database software, and streaming video content software.

(56) In at least one embodiment, application(s) **742** included in application layer **740** may include one or more types of applications used by at least portions of node C.R.s **716(1)-716(N)**, grouped computing resources **714**, and/or distributed file system **728** of framework layer **720**. In at least one embodiment, one or more types of applications may include, but are not limited to, any number of a genomics application, a cognitive compute, application and a machine learning application, including training or inferencing software, machine learning framework software (e.g., PyTorch, TensorFlow, Caffe, etc.) or other machine learning applications used in conjunction with one or more embodiments.

(57) In at least one embodiment, any of configuration manager **724**, resource manager **726**, and resource orchestrator **712** may implement any number and type of self-modifying actions based on any amount and type of data acquired in any technically feasible fashion. In at least one embodiment, self-modifying actions may relieve a data center operator of data center **700** from making possibly bad configuration decisions and possibly avoiding underutilized and/or poor performing portions of a data center.

(58) In at least one embodiment, data center **700** may include tools, services, software or other

resources to train one or more machine learning models or predict or infer information using one or more machine learning models according to one or more embodiments described herein. For example, in at least one embodiment, a machine learning model may be trained by calculating weight parameters according to a neural network architecture using software and computing resources described above with respect to data center **700**. In at least one embodiment, trained machine learning models corresponding to one or more neural networks may be used to infer or predict information using resources described above with respect to data center **700** by using weight parameters calculated through one or more training techniques described herein.

(59) In at least one embodiment, data center may use CPUs, application-specific integrated circuits (ASICs), GPUs, FPGAs, or other hardware to perform training and/or inferencing using above-described resources. Moreover, one or more software and/or hardware resources described above may be configured as a service to allow users to train or performing inferencing of information, such as image recognition, speech recognition, or other artificial intelligence services.

(60) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in data center **700** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(61) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(62) Computer Systems

(63) FIG. **8** is a block diagram illustrating an exemplary computer system, which may be a system with interconnected devices and components, a system-on-a-chip (SOC) or some combination thereof formed with a processor that may include execution units to execute an instruction, according to at least one embodiment. In at least one embodiment, a computer system **800** may include, without limitation, a component, such as a processor **802** to employ execution units including logic to perform algorithms for process data, in accordance with present disclosure, such as in embodiment described herein. In at least one embodiment, computer system **800** may include processors, such as PENTIUM® Processor family, Xeon™, Itanium®, XScale™ and/or StrongARM™, Intel® Core™, or Intel® Nervana™ microprocessors available from Intel Corporation of Santa Clara, California, although other systems (including PCs having other microprocessors, engineering workstations, set-top boxes and like) may also be used. In at least one embodiment, computer system **800** may execute a version of WINDOWS operating system available from Microsoft Corporation of Redmond, Wash., although other operating systems (UNIX and Linux, for example), embedded software, and/or graphical user interfaces, may also be used.

(64) Embodiments may be used in other devices such as handheld devices and embedded applications. Some examples of handheld devices include cellular phones, Internet Protocol devices, digital cameras, personal digital assistants (“PDAs”), and handheld PCs. In at least one embodiment, embedded applications may include a microcontroller, a digital signal processor (“DSP”), system on a chip, network computers (“NetPCs”), set-top boxes, network hubs, wide area network (“WAN”) switches, or any other system that may perform one or more instructions in accordance with at least one embodiment.

(65) In at least one embodiment, computer system **800** may include, without limitation, processor **802** that may include, without limitation, one or more execution units **808** to perform machine learning model training and/or inferencing according to techniques described herein. In at least one embodiment, computer system **800** is a single processor desktop or server system, but in another embodiment, computer system **800** may be a multiprocessor system. In at least one embodiment, processor **802** may include, without limitation, a complex instruction set computer (“CISC”) microprocessor, a reduced instruction set computing (“RISC”) microprocessor, a very long

instruction word (“VLIW”) microprocessor, a processor implementing a combination of instruction sets, or any other processor device, such as a digital signal processor, for example. In at least one embodiment, processor **802** may be coupled to a processor bus **810** that may transmit data signals between processor **802** and other components in computer system **800**.

(66) In at least one embodiment, processor **802** may include, without limitation, a Level 1 (“L1”) internal cache memory (“cache”) **804**. In at least one embodiment, processor **802** may have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory may reside external to processor **802**. Other embodiments may also include a combination of both internal and external caches depending on particular implementation and needs. In at least one embodiment, a register file **806** may store different types of data in various registers including, without limitation, integer registers, floating point registers, status registers, and an instruction pointer register.

(67) In at least one embodiment, execution unit **808**, including, without limitation, logic to perform integer and floating point operations, also resides in processor **802**. In at least one embodiment, processor **802** may also include a microcode (“ucode”) read only memory (“ROM”) that stores microcode for certain macro instructions. In at least one embodiment, execution unit **808** may include logic to handle a packed instruction set **809**. In at least one embodiment, by including packed instruction set **809** in an instruction set of a general-purpose processor, along with associated circuitry to execute instructions, operations used by many multimedia applications may be performed using packed data in processor **802**. In at least one embodiment, many multimedia applications may be accelerated and executed more efficiently by using a full width of a processor's data bus for performing operations on packed data, which may eliminate a need to transfer smaller units of data across that processor's data bus to perform one or more operations one data element at a time.

(68) In at least one embodiment, execution unit **808** may also be used in microcontrollers, embedded processors, graphics devices, DSPs, and other types of logic circuits. In at least one embodiment, computer system **800** may include, without limitation, a memory **820**. In at least one embodiment, memory **820** may be a Dynamic Random Access Memory (“DRAM”) device, a Static Random Access Memory (“SRAM”) device, a flash memory device, or another memory device. In at least one embodiment, memory **820** may store instruction(s) **819** and/or data **821** represented by data signals that may be executed by processor **802**.

(69) In at least one embodiment, a system logic chip may be coupled to processor bus **810** and memory **820**. In at least one embodiment, a system logic chip may include, without limitation, a memory controller hub (“MCH”) **816**, and processor **802** may communicate with MCH **816** via processor bus **810**. In at least one embodiment, MCH **816** may provide a high bandwidth memory path **818** to memory **820** for instruction and data storage and for storage of graphics commands, data and textures. In at least one embodiment, MCH **816** may direct data signals between processor **802**, memory **820**, and other components in computer system **800** and to bridge data signals between processor bus **810**, memory **820**, and a system I/O interface **822**. In at least one embodiment, a system logic chip may provide a graphics port for coupling to a graphics controller. In at least one embodiment, MCH **816** may be coupled to memory **820** through high bandwidth memory path **818** and a graphics/video card **812** may be coupled to MCH **816** through an Accelerated Graphics Port (“AGP”) interconnect **814**.

(70) In at least one embodiment, computer system **800** may use system I/O interface **822** as a proprietary hub interface bus to couple MCH **816** to an I/O controller hub (“ICH”) **830**. In at least one embodiment, ICH **830** may provide direct connections to some I/O devices via a local I/O bus. In at least one embodiment, a local I/O bus may include, without limitation, a high-speed I/O bus for connecting peripherals to memory **820**, a chipset, and processor **802**. Examples may include, without limitation, an audio controller **829**, a firmware hub (“flash BIOS”) **828**, a wireless transceiver **826**, a data storage **824**, a legacy I/O controller **823** containing user input and keyboard

interfaces **825**, a serial expansion port **827**, such as a Universal Serial Bus (“USB”) port, and a network controller **834**. In at least one embodiment, data storage **824** may comprise a hard disk drive, a floppy disk drive, a CD-ROM device, a flash memory device, or other mass storage device. (71) In at least one embodiment, FIG. **8** illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG. **8** may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. **8** may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of computer system **800** are interconnected using compute express link (CXL) interconnects.

(72) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in computer system **800** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(73) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(74) FIG. **9** is a block diagram illustrating an electronic device **900** for utilizing a processor **910**, according to at least one embodiment. In at least one embodiment, electronic device **900** may be, for example and without limitation, a notebook, a tower server, a rack server, a blade server, a laptop, a desktop, a tablet, a mobile device, a phone, an embedded computer, or any other suitable electronic device.

(75) In at least one embodiment, electronic device **900** may include, without limitation, processor **910** communicatively coupled to any suitable number or kind of components, peripherals, modules, or devices. In at least one embodiment, processor **910** is coupled using a bus or interface, such as a I.sup.2C bus, a System Management Bus (“SMBus”), a Low Pin Count (LPC) bus, a Serial Peripheral Interface (“SPI”), a High Definition Audio (“HDA”) bus, a Serial Advance Technology Attachment (“SATA”) bus, a Universal Serial Bus (“USB”) (versions 1, 2, 3, etc.), or a Universal Asynchronous Receiver/Transmitter (“UART”) bus. In at least one embodiment, FIG. **9** illustrates a system, which includes interconnected hardware devices or “chips”, whereas in other embodiments, FIG. **9** may illustrate an exemplary SoC. In at least one embodiment, devices illustrated in FIG. **9** may be interconnected with proprietary interconnects, standardized interconnects (e.g., PCIe) or some combination thereof. In at least one embodiment, one or more components of FIG. **9** are interconnected using compute express link (CXL) interconnects.

(76) In at least one embodiment, FIG. **9** may include a display **924**, a touch screen **925**, a touch pad **930**, a Near Field Communications unit (“NFC”) **945**, a sensor hub **940**, a thermal sensor **946**, an Express Chipset (“EC”) **935**, a Trusted Platform Module (“TPM”) **938**, BIOS/firmware/flash memory (“BIOS, FW Flash”) **922**, a DSP **960**, a drive **920** such as a Solid State Disk (“SSD”) or a Hard Disk Drive (“HDD”), a wireless local area network unit (“WLAN”) **950**, a Bluetooth unit **952**, a Wireless Wide Area Network unit (“WWAN”) **956**, a Global Positioning System (GPS) unit **955**, a camera (“USB 3.0 camera”) **954** such as a USB 3.0 camera, and/or a Low Power Double Data Rate (“LPDDR”) memory unit (“LPDDR3”) **915** implemented in, for example, an LPDDR3 standard. These components may each be implemented in any suitable manner.

(77) In at least one embodiment, other components may be communicatively coupled to processor **910** through components described herein. In at least one embodiment, an accelerometer **941**, an ambient light sensor (“ALS”) **942**, a compass **943**, and a gyroscope **944** may be communicatively coupled to sensor hub **940**. In at least one embodiment, a thermal sensor **939**, a fan **937**, a keyboard **936**, and touch pad **930** may be communicatively coupled to EC **935**. In at least one embodiment, speakers **963**, headphones **964**, and a microphone (“mic”) **965** may be communicatively coupled to an audio unit (“audio codec and class D amp”) **962**, which may in turn be communicatively coupled to DSP **960**. In at least one embodiment, audio unit **962** may include, for example and

without limitation, an audio coder/decoder (“codec”) and a class D amplifier. In at least one embodiment, a SIM card (“SIM”) **957** may be communicatively coupled to WWAN unit **956**. In at least one embodiment, components such as WLAN unit **950** and Bluetooth unit **952**, as well as WWAN unit **956** may be implemented in a Next Generation Form Factor (“NGFF”).

(78) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in electronic device **900** or inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(79) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(80) FIG. **10** illustrates a computer system **1000**, according to at least one embodiment. In at least one embodiment, computer system **1000** is configured to implement various processes and methods described throughout this disclosure.

(81) In at least one embodiment, computer system **1000** comprises, without limitation, at least one central processing unit (“CPU”) **1002** that is connected to a communication bus **1010** implemented using any suitable protocol, such as PCI (“Peripheral Component Interconnect”), peripheral component interconnect express (“PCI-Express”), AGP (“Accelerated Graphics Port”), HyperTransport, or any other bus or point-to-point communication protocol(s). In at least one embodiment, computer system **1000** includes, without limitation, a main memory **1004** and control logic (e.g., implemented as hardware, software, or a combination thereof) and data are stored in main memory **1004**, which may take form of random access memory (“RAM”). In at least one embodiment, a network interface subsystem (“network interface”) **1022** provides an interface to other computing devices and networks for receiving data from and transmitting data to other systems with computer system **1000**.

(82) In at least one embodiment, computer system **1000**, in at least one embodiment, includes, without limitation, input devices **1008**, a parallel processing system **1012**, and display devices **1006** that can be implemented using a conventional cathode ray tube (“CRT”), a liquid crystal display (“LCD”), a light emitting diode (“LED”) display, a plasma display, or other suitable display technologies. In at least one embodiment, user input is received from input devices **1008** such as keyboard, mouse, touchpad, microphone, etc. In at least one embodiment, each module described herein can be situated on a single semiconductor platform to form a processing system.

(83) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in computer system **1000** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(84) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(85) FIG. **11** illustrates a computer system **1100**, according to at least one embodiment. In at least one embodiment, computer system **1100** includes, without limitation, a computer **1110** and a USB stick **1120**. In at least one embodiment, computer **1110** may include, without limitation, any number and type of processor(s) (not shown) and a memory (not shown). In at least one embodiment, computer **1110** includes, without limitation, a server, a cloud instance, a laptop, and a desktop computer.

(86) In at least one embodiment, USB stick **1120** includes, without limitation, a processing unit **1130**, a USB interface **1140**, and USB interface logic **1150**. In at least one embodiment, processing unit **1130** may be any instruction execution system, apparatus, or device capable of executing instructions. In at least one embodiment, processing unit **1130** may include, without limitation, any number and type of processing cores (not shown). In at least one embodiment, processing unit **1130**

comprises an application specific integrated circuit (“ASIC”) that is optimized to perform any amount and type of operations associated with machine learning. For instance, in at least one embodiment, processing unit **1130** is a tensor processing unit (“TPC”) that is optimized to perform machine learning inference operations. In at least one embodiment, processing unit **1130** is a vision processing unit (“VPU”) that is optimized to perform machine vision and machine learning inference operations.

(87) In at least one embodiment, USB interface **1140** may be any type of USB connector or USB socket. For instance, in at least one embodiment, USB interface **1140** is a USB 3.0 Type-C socket for data and power. In at least one embodiment, USB interface **1140** is a USB 3.0 Type-A connector. In at least one embodiment, USB interface logic **1150** may include any amount and type of logic that enables processing unit **1130** to interface with devices (e.g., computer **1110**) via USB connector **1140**.

(88) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in computer system **1100** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(89) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(90) FIG. **12A** illustrates an exemplary architecture in which a plurality of GPUs **1210(1)-1210(N)** is communicatively coupled to a plurality of multi-core processors **1205(1)-1205(M)** over high-speed links **1240(1)-1240(N)** (e.g., buses, point-to-point interconnects, etc.). In at least one embodiment, high-speed links **1240(1)-1240(N)** support a communication throughput of 4 GB/s, 30 GB/s, 80 GB/s or higher. In at least one embodiment, various interconnect protocols may be used including, but not limited to, PCIe 4.0 or 5.0 and NVLink 2.0. In various figures, “N” and “M” represent positive integers, values of which may be different from figure to figure. In at least one embodiment, one or more GPUs in a plurality of GPUs **1210(1)-1210(N)** includes one or more graphics cores (also referred to simply as “cores”) **1500** as disclosed in FIGS. **15A** and **15B**. In at least one embodiment, one or more graphics cores **1500** may be referred to as streaming multiprocessors (“SMs”), stream processors (“SPs”), stream processing units (“SPUs”), compute units (“CUs”), execution units (“EUs”), and/or slices, where a slice in this context can refer to a portion of processing resources in a processing unit (e.g., 16 cores, a ray tracing unit, a thread director or scheduler).

(91) In addition, and in at least one embodiment, two or more of GPUs **1210** are interconnected over high-speed links **1229(1)-1229(2)**, which may be implemented using similar or different protocols/links than those used for high-speed links **1240(1)-1240(N)**. Similarly, two or more of multi-core processors **1205** may be connected over a high-speed link **1228** which may be symmetric multi-processor (SMP) buses operating at 20 GB/s, 30 GB/s, 120 GB/s or higher. Alternatively, all communication between various system components shown in FIG. **12A** may be accomplished using similar protocols/links (e.g., over a common interconnection fabric).

(92) In at least one embodiment, each multi-core processor **1205** is communicatively coupled to a processor memory **1201(1)-1201(M)**, via memory interconnects **1226(1)-1226(M)**, respectively, and each GPU **1210(1)-1210(N)** is communicatively coupled to GPU memory **1220(1)-1220(N)** over GPU memory interconnects **1250(1)-1250(N)**, respectively. In at least one embodiment, memory interconnects **1226** and **1250** may utilize similar or different memory access technologies. By way of example, and not limitation, processor memories **1201(1)-1201(M)** and GPU memories **1220** may be volatile memories such as dynamic random access memories (DRAMs) (including stacked DRAMs), Graphics DDR SDRAM (GDDR) (e.g., GDDR5, GDDR6), or High Bandwidth Memory (HBM) and/or may be non-volatile memories such as 3D XPoint or Nano-Ram. In at least one embodiment, some portion of processor memories **1201** may be volatile memory and another

portion may be non-volatile memory (e.g., using a two-level memory (2LM) hierarchy).

(93) As described herein, although various multi-core processors **1205** and GPUs **1210** may be physically coupled to a particular memory **1201**, **1220**, respectively, and/or a unified memory architecture may be implemented in which a virtual system address space (also referred to as “effective address” space) is distributed among various physical memories. For example, processor memories **1201(1)-1201(M)** may each comprise 64 GB of system memory address space and GPU memories **1220(1)-1220(N)** may each comprise 32 GB of system memory address space resulting in a total of 256 GB addressable memory when M=2 and N=4. Other values for N and M are possible.

(94) FIG. **12B** illustrates additional details for an interconnection between a multi-core processor **1207** and a graphics acceleration module **1246** in accordance with one exemplary embodiment. In at least one embodiment, graphics acceleration module **1246** may include one or more GPU chips integrated on a line card which is coupled to processor **1207** via high-speed link **1240** (e.g., a PCIe bus, NVLink, etc.). In at least one embodiment, graphics acceleration module **1246** may alternatively be integrated on a package or chip with processor **1207**.

(95) In at least one embodiment, processor **1207** includes a plurality of cores **1260A-1260D** (which may be referred to as “execution units”), each with a translation lookaside buffer (“TLB”) **1261A-1261D** and one or more caches **1262A-1262D**. In at least one embodiment, cores **1260A-1260D** may include various other components for executing instructions and processing data that are not illustrated. In at least one embodiment, caches **1262A-1262D** may comprise Level 1 (L1) and Level 2 (L2) caches. In addition, one or more shared caches **1256** may be included in caches **1262A-1262D** and shared by sets of cores **1260A-1260D**. For example, one embodiment of processor **1207** includes 24 cores, each with its own L1 cache, twelve shared L2 caches, and twelve shared L3 caches. In this embodiment, one or more L2 and L3 caches are shared by two adjacent cores. In at least one embodiment, processor **1207** and graphics acceleration module **1246** connect with system memory **1214**, which may include processor memories **1201(1)-1201(M)** of FIG. **12A**.

(96) In at least one embodiment, coherency is maintained for data and instructions stored in various caches **1262A-1262D**, **1256** and system memory **1214** via inter-core communication over a coherence bus **1264**. In at least one embodiment, for example, each cache may have cache coherency logic/circuitry associated therewith to communicate to over coherence bus **1264** in response to detected reads or writes to particular cache lines. In at least one embodiment, a cache snooping protocol is implemented over coherence bus **1264** to snoop cache accesses.

(97) In at least one embodiment, a proxy circuit **1225** communicatively couples graphics acceleration module **1246** to coherence bus **1264**, allowing graphics acceleration module **1246** to participate in a cache coherence protocol as a peer of cores **1260A-1260D**. In particular, in at least one embodiment, an interface **1235** provides connectivity to proxy circuit **1225** over high-speed link **1240** and an interface **1237** connects graphics acceleration module **1246** to high-speed link **1240**.

(98) In at least one embodiment, an accelerator integration circuit **1236** provides cache management, memory access, context management, and interrupt management services on behalf of a plurality of graphics processing engines **1231(1)-1231(N)** of graphics acceleration module **1246**. In at least one embodiment, graphics processing engines **1231(1)-1231(N)** may each comprise a separate graphics processing unit (GPU). In at least one embodiment, plurality of graphics processing engines **1231(1)-1231(N)** of graphics acceleration module **1246** include one or more graphics cores **1500** as discussed in connection with FIGS. **15A** and **15B**. In at least one embodiment, graphics processing engines **1231(1)-1231(N)** alternatively may comprise different types of graphics processing engines within a GPU, such as graphics execution units, media processing engines (e.g., video encoders/decoders), samplers, and blit engines. In at least one embodiment, graphics acceleration module **1246** may be a GPU with a plurality of graphics processing engines **1231(1)-1231(N)** or graphics processing engines **1231(1)-1231(N)** may be

individual GPUs integrated on a common package, line card, or chip.

(99) In at least one embodiment, accelerator integration circuit **1236** includes a memory management unit (MMU) **1239** for performing various memory management functions such as virtual-to-physical memory translations (also referred to as effective-to-real memory translations) and memory access protocols for accessing system memory **1214**. In at least one embodiment, MMU **1239** may also include a translation lookaside buffer (TLB) (not shown) for caching virtual/effective to physical/real address translations. In at least one embodiment, a cache **1238** can store commands and data for efficient access by graphics processing engines **1231(1)-1231(N)**. In at least one embodiment, data stored in cache **1238** and graphics memories **1233(1)-1233(M)** is kept coherent with core caches **1262A-1262D**, **1256** and system memory **1214**, possibly using a fetch unit **1244**. As mentioned, this may be accomplished via proxy circuit **1225** on behalf of cache **1238** and memories **1233(1)-1233(M)** (e.g., sending updates to cache **1238** related to modifications/accesses of cache lines on processor caches **1262A-1262D**, **1256** and receiving updates from cache **1238**).

(100) In at least one embodiment, a set of registers **1245** store context data for threads executed by graphics processing engines **1231(1)-1231(N)** and a context management circuit **1248** manages thread contexts. For example, context management circuit **1248** may perform save and restore operations to save and restore contexts of various threads during contexts switches (e.g., where a first thread is saved and a second thread is stored so that a second thread can be execute by a graphics processing engine). For example, on a context switch, context management circuit **1248** may store current register values to a designated region in memory (e.g., identified by a context pointer). It may then restore register values when returning to a context. In at least one embodiment, an interrupt management circuit **1247** receives and processes interrupts received from system devices.

(101) In at least one embodiment, virtual/effective addresses from a graphics processing engine **1231** are translated to real/physical addresses in system memory **1214** by MMU **1239**. In at least one embodiment, accelerator integration circuit **1236** supports multiple (e.g., 4, 8, 16) graphics accelerator modules **1246** and/or other accelerator devices. In at least one embodiment, graphics accelerator module **1246** may be dedicated to a single application executed on processor **1207** or may be shared between multiple applications. In at least one embodiment, a virtualized graphics execution environment is presented in which resources of graphics processing engines **1231(1)-1231(N)** are shared with multiple applications or virtual machines (VMs). In at least one embodiment, resources may be subdivided into “slices” which are allocated to different VMs and/or applications based on processing requirements and priorities associated with VMs and/or applications.

(102) In at least one embodiment, accelerator integration circuit **1236** performs as a bridge to a system for graphics acceleration module **1246** and provides address translation and system memory cache services. In addition, in at least one embodiment, accelerator integration circuit **1236** may provide virtualization facilities for a host processor to manage virtualization of graphics processing engines **1231(1)-1231(N)**, interrupts, and memory management.

(103) In at least one embodiment, because hardware resources of graphics processing engines **1231(1)-1231(N)** are mapped explicitly to a real address space seen by host processor **1207**, any host processor can address these resources directly using an effective address value. In at least one embodiment, one function of accelerator integration circuit **1236** is physical separation of graphics processing engines **1231(1)-1231(N)** so that they appear to a system as independent units.

(104) In at least one embodiment, one or more graphics memories **1233(1)-1233(M)** are coupled to each of graphics processing engines **1231(1)-1231(N)**, respectively and $N=M$. In at least one embodiment, graphics memories **1233(1)-1233(M)** store instructions and data being processed by each of graphics processing engines **1231(1)-1231(N)**. In at least one embodiment, graphics memories **1233(1)-1233(M)** may be volatile memories such as DRAMs (including stacked

DRAMs), GDDR memory (e.g., GDDR5, GDDR6), or HBM, and/or may be non-volatile memories such as 3D XPoint or Nano-Ram.

(105) In at least one embodiment, to reduce data traffic over high-speed link **1240**, biasing techniques can be used to ensure that data stored in graphics memories **1233(1)-1233(M)** is data that will be used most frequently by graphics processing engines **1231(1)-1231(N)** and preferably not used by cores **1260A-1260D** (at least not frequently). Similarly, in at least one embodiment, a biasing mechanism attempts to keep data needed by cores (and preferably not graphics processing engines **1231(1)-1231(N)**) within caches **1262A-1262D**, **1256** and system memory **1214**.

(106) FIG. **13** illustrates exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

(107) FIG. **13** is a block diagram illustrating an exemplary system on a chip integrated circuit **1300** that may be fabricated using one or more IP cores, according to at least one embodiment. In at least one embodiment, integrated circuit **1300** includes one or more application processor(s) **1305** (e.g., CPUs), at least one graphics processor **1310**, and may additionally include an image processor **1315** and/or a video processor **1320**, any of which may be a modular IP core. In at least one embodiment, integrated circuit **1300** includes peripheral or bus logic including a USB controller **1325**, a UART controller **1330**, an SPI/SDIO controller **1335**, and an I.sup.2S/I.sup.2C controller **1340**. In at least one embodiment, integrated circuit **1300** can include a display device **1345** coupled to one or more of a high-definition multimedia interface (HDMI) controller **1350** and a mobile industry processor interface (MIPI) display interface **1355**. In at least one embodiment, storage may be provided by a flash memory subsystem **1360** including flash memory and a flash memory controller. In at least one embodiment, a memory interface may be provided via a memory controller **1365** for access to SDRAM or SRAM memory devices. In at least one embodiment, some integrated circuits additionally include an embedded security engine **1370**.

(108) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in integrated circuit **1300** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(109) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(110) FIGS. **14A-14B** illustrate exemplary integrated circuits and associated graphics processors that may be fabricated using one or more IP cores, according to various embodiments described herein. In addition to what is illustrated, other logic and circuits may be included in at least one embodiment, including additional graphics processors/cores, peripheral interface controllers, or general-purpose processor cores.

(111) FIGS. **14A-14B** are block diagrams illustrating exemplary graphics processors for use within an SoC, according to embodiments described herein. FIG. **14A** illustrates an exemplary graphics processor **1410** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to at least one embodiment. FIG. **14B** illustrates an additional exemplary graphics processor **1440** of a system on a chip integrated circuit that may be fabricated using one or more IP cores, according to at least one embodiment. In at least one embodiment, graphics processor **1410** of FIG. **14A** is a low power graphics processor core. In at least one embodiment, graphics processor **1440** of FIG. **14B** is a higher performance graphics processor core. In at least one embodiment, each of graphics processors **1410**, **1440** can be variants of graphics processor **1310** of FIG. **13**.

(112) In at least one embodiment, graphics processor **1410** includes a vertex processor **1405** and

one or more fragment processor(s) **1415A-1415N** (e.g., **1415A**, **1415B**, **1415C**, **1415D**, through **1415N-1**, and **1415N**). In at least one embodiment, graphics processor **1410** can execute different shader programs via separate logic, such that vertex processor **1405** is optimized to execute operations for vertex shader programs, while one or more fragment processor(s) **1415A-1415N** execute fragment (e.g., pixel) shading operations for fragment or pixel shader programs. In at least one embodiment, vertex processor **1405** performs a vertex processing stage of a 3D graphics pipeline and generates primitives and vertex data. In at least one embodiment, fragment processor(s) **1415A-1415N** use primitive and vertex data generated by vertex processor **1405** to produce a framebuffer that is displayed on a display device. In at least one embodiment, fragment processor(s) **1415A-1415N** are optimized to execute fragment shader programs as provided for in an OpenGL API, which may be used to perform similar operations as a pixel shader program as provided for in a Direct 3D API.

(113) In at least one embodiment, graphics processor **1410** additionally includes one or more memory management units (MMUs) **1420A-1420B**, cache(s) **1425A-1425B**, and circuit interconnect(s) **1430A-1430B**. In at least one embodiment, one or more MMU(s) **1420A-1420B** provide for virtual to physical address mapping for graphics processor **1410**, including for vertex processor **1405** and/or fragment processor(s) **1415A-1415N**, which may reference vertex or image/texture data stored in memory, in addition to vertex or image/texture data stored in one or more cache(s) **1425A-1425B**. In at least one embodiment, one or more MMU(s) **1420A-1420B** may be synchronized with other MMUs within a system, including one or more MMUs associated with one or more application processor(s) **1305**, image processors **1315**, and/or video processors **1320** of FIG. **13**, such that each processor **1305-1320** can participate in a shared or unified virtual memory system. In at least one embodiment, one or more circuit interconnect(s) **1430A-1430B** enable graphics processor **1410** to interface with other IP cores within SoC, either via an internal bus of SoC or via a direct connection.

(114) In at least one embodiment, graphics processor **1440** includes one or more shader core(s) **1455A-1455N** (e.g., **1455A**, **1455B**, **1455C**, **1455D**, **1455E**, **1455F**, through **1455N-1**, and **1455N**) as shown in FIG. **14B**, which provides for a unified shader core architecture in which a single core or type or core can execute all types of programmable shader code, including shader program code to implement vertex shaders, fragment shaders, and/or compute shaders. In at least one embodiment, a number of shader cores can vary. In at least one embodiment, graphics processor **1440** includes an inter-core task manager **1445**, which acts as a thread dispatcher to dispatch execution threads to one or more shader cores **1455A-1455N** and a tiling unit **1458** to accelerate tiling operations for tile-based rendering, in which rendering operations for a scene are subdivided in image space, for example to exploit local spatial coherence within a scene or to optimize use of internal caches.

(115) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in graphics processors **1410** and/or **1140** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(116) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(117) FIGS. **15A-15B** illustrate additional exemplary graphics processor logic according to embodiments described herein. FIG. **15A** illustrates a graphics core **1500** that may be included within graphics processor **1310** of FIG. **13**, in at least one embodiment, and may be a unified shader core **1455A-1455N** as in FIG. **14B** in at least one embodiment. FIG. **15B** illustrates a highly-parallel general-purpose graphics processing unit (“GPGPU”) **1530** suitable for deployment on a multi-chip module in at least one embodiment.

(118) In at least one embodiment, graphics core **1500** includes a shared instruction cache **1502**, a texture unit **1518**, and a cache/shared memory **1520** (e.g., including L1, L2, L3, last level cache, or other caches) that are common to execution resources within graphics core **1500**. In at least one embodiment, graphics core **1500** can include multiple slices **1501A-1501N** or a partition for each core, and a graphics processor can include multiple instances of graphics core **1500**. In at least one embodiment, each slice **1501A-1501N** refers to graphics core **1500**. In at least one embodiment, slices **1501A-1501N** have sub-slices, which are part of a slice **1501A-1501N**. In at least one embodiment, slices **1501A-1501N** are independent of other slices or dependent on other slices. In at least one embodiment, slices **1501A-1501N** can include support logic including a local instruction cache **1504A-1504N**, a thread scheduler (sequencer) **1506A-1506N**, a thread dispatcher **1508A-1508N**, and a set of registers **1510A-1510N**. In at least one embodiment, slices **1501A-1501N** can include a set of additional function units (AFUs **1512A-1512N**), floating-point units (FPUs **1514A-1514N**), integer arithmetic logic units (ALUs **1516A-1516N**), address computational units (ACUs **1513A-1513N**), double-precision floating-point units (DPFPUs **1515A-1515N**), and matrix processing units (MPUs **1517A-1517N**).

(119) In at least one embodiment, each slice **1501A-1501N** includes one or more engines for floating point and integer vector operations and one or more engines to accelerate convolution and matrix operations in AI, machine learning, or large dataset workloads. In at least one embodiment, one or more slices **1501A-1501N** include one or more vector engines to compute a vector (e.g., compute mathematical operations for vectors). In at least one embodiment, a vector engine can compute a vector operation in 16-bit floating point (also referred to as “FP16”), 32-bit floating point (also referred to as “FP32”), or 64-bit floating point (also referred to as “FP64”). In at least one embodiment, one or more slices **1501A-1501N** includes 16 vector engines that are paired with 16 matrix math units to compute matrix/tensor operations, where vector engines and math units are exposed via matrix extensions. In at least one embodiment, a slice a specified portion of processing resources of a processing unit, e.g., 16 cores and a ray tracing unit or 8 cores, a thread scheduler, a thread dispatcher, and additional functional units for a processor. In at least one embodiment, graphics core **1500** includes one or more matrix engines to compute matrix operations, e.g., when computing tensor operations.

(120) In at least one embodiment, one or more slices **1501A-1501N** includes one or more ray tracing units to compute ray tracing operations (e.g., 16 ray tracing units per slice slices **1501A-1501N**). In at least one embodiment, a ray tracing unit computes ray traversal, triangle intersection, bounding box intersect, or other ray tracing operations.

(121) In at least one embodiment, one or more slices **1501A-1501N** includes a media slice that encodes, decodes, and/or transcodes data; scales and/or format converts data; and/or performs video quality operations on video data.

(122) In at least one embodiment, one or more slices **1501A-1501N** are linked to L2 cache and memory fabric, link connectors, high-bandwidth memory (HBM) (e.g., HBM2e, HBM3) stacks, and a media engine. In at least one embodiment, one or more slices **1501A-1501N** include multiple cores (e.g., 16 cores) and multiple ray tracing units (e.g., 16) paired to each core. In at least one embodiment, one or more slices **1501A-1501N** has one or more L1 caches. In at least one embodiment, one or more slices **1501A-1501N** include one or more vector engines; one or more instruction caches to store instructions; one or more L1 caches to cache data; one or more shared local memories (SLMs) to store data, e.g., corresponding to instructions; one or more samplers to sample data; one or more ray tracing units to perform ray tracing operations; one or more geometries to perform operations in geometry pipelines and/or apply geometric transformations to vertices or polygons; one or more rasterizers to describe an image in vector graphics format (e.g., shape) and convert it into a raster image (e.g., a series of pixels, dots, or lines, which when displayed together, create an image that is represented by shapes); one or more a Hierarchical Depth Buffer (Hiz) to buffer data; and/or one or more pixel backends. In at least one embodiment, a

slice **1501A-1501N** includes a memory fabric, e.g., an L2 cache.

(123) In at least one embodiment, FPUs **1514A-1514N** can perform single-precision (32-bit) and half-precision (16-bit) floating point operations, while DPFPUs **1515A-1515N** perform double precision (64-bit) floating point operations. In at least one embodiment, ALUs **1516A-1516N** can perform variable precision integer operations at 8-bit, 16-bit, and 32-bit precision, and can be configured for mixed precision operations. In at least one embodiment, MPUs **1517A-1517N** can also be configured for mixed precision matrix operations, including half-precision floating point and 8-bit integer operations. In at least one embodiment, MPUs **1517-1517N** can perform a variety of matrix operations to accelerate machine learning application frameworks, including enabling support for accelerated general matrix to matrix multiplication (GEMM). In at least one embodiment, AFUs **1512A-1512N** can perform additional logic operations not supported by floating-point or integer units, including trigonometric operations (e.g., sine, cosine). Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in graphics core **1500** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(124) In at least one embodiment, graphics core **1500** includes an interconnect and a link fabric sublayer that is attached to a switch and a GPU-GPU bridge that enables multiple graphics processors **1500** (e.g., 8) to be interlinked without glue to each other with load/store units (LSUs), data transfer units, and sync semantics across multiple graphics processors **1500**. In at least one embodiment, interconnects include standardized interconnects (e.g., PCIe) or some combination thereof.

(125) In at least one embodiment, graphics core **1500** includes multiple tiles. In at least one embodiment, a tile is an individual die or one or more dies, where individual dies can be connected with an interconnect (e.g., embedded multi-die interconnect bridge (EMIB)). In at least one embodiment, graphics core **1500** includes a compute tile, a memory tile (e.g., where a memory tile can be exclusively accessed by different tiles or different chipsets such as a Rambo tile), substrate tile, a base tile, a HMB tile, a link tile, and EMIB tile, where all tiles are packaged together in graphics core **1500** as part of a GPU. In at least one embodiment, graphics core **1500** can include multiple tiles in a single package (also referred to as a “multi tile package”). In at least one embodiment, a compute tile can have 8 graphics cores **1500**, an L1 cache; and a base tile can have a host interface with PCIe 5.0, HBM2e, MDFI, and EMIB, a link tile with 8 links, 8 ports with an embedded switch. In at least one embodiment, tiles are connected with face-to-face (F2F) chip-on-chip bonding through fine-pitched, 36-micron, microbumps (e.g., copper pillars). In at least one embodiment, graphics core **1500** includes memory fabric, which includes memory, and is tile that is accessible by multiple tiles. In at least one embodiment, graphics core **1500** stores, accesses, or loads its own hardware contexts in memory, where a hardware context is a set of data loaded from registers before a process resumes, and where a hardware context can indicate a state of hardware (e.g., state of a GPU).

(126) In at least one embodiment, graphics core **1500** includes serializer/deserializer (SERDES) circuitry that converts a serial data stream to a parallel data stream, or converts a parallel data stream to a serial data stream.

(127) In at least one embodiment, graphics core **1500** includes a high speed coherent unified fabric (GPU to GPU), load/store units, bulk data transfer and sync semantics, and connected GPUs through an embedded switch, where a GPU-GPU bridge is controlled by a controller.

(128) In at least one embodiment, graphics core **1500** performs an API, where said API abstracts hardware of graphics core **1500** and access libraries with instructions to perform math operations (e.g., math kernel library), deep neural network operations (e.g., deep neural network library), vector operations, collective communications, thread building blocks, video processing, data

analytics library, and/or ray tracing operations.

(129) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(130) FIG. 15B illustrates a general-purpose processing unit (GPGPU) 1530 that can be configured to enable highly-parallel compute operations to be performed by an array of graphics processing units, in at least one embodiment. In at least one embodiment, GPGPU 1530 can be linked directly to other instances of GPGPU 1530 to create a multi-GPU cluster to improve training speed for deep neural networks. In at least one embodiment, GPGPU 1530 includes a host interface 1532 to enable a connection with a host processor. In at least one embodiment, host interface 1532 is a PCI Express interface. In at least one embodiment, host interface 1532 can be a vendor-specific communications interface or communications fabric. In at least one embodiment, GPGPU 1530 receives commands from a host processor and uses a global scheduler 1534 (which may be referred to as a thread sequencer and/or asynchronous compute engine) to distribute execution threads associated with those commands to a set of compute clusters 1536A-1536H. In at least one embodiment, compute clusters 1536A-1536H share a cache memory 1538. In at least one embodiment, cache memory 1538 can serve as a higher-level cache for cache memories within compute clusters 1536A-1536H.

(131) In at least one embodiment, GPGPU 1530 includes memory 1544A-1544B coupled with compute clusters 1536A-1536H via a set of memory controllers 1542A-1542B (e.g., one or more controllers for HBM2e). In at least one embodiment, memory 1544A-1544B can include various types of memory devices including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory.

(132) In at least one embodiment, compute clusters 1536A-1536H each include a set of graphics cores, such as graphics core 1500 of FIG. 15A, which can include multiple types of integer and floating point logic units that can perform computational operations at a range of precisions including suited for machine learning computations. For example, in at least one embodiment, at least a subset of floating point units in each of compute clusters 1536A-1536H can be configured to perform 16-bit or 32-bit floating point operations, while a different subset of floating point units can be configured to perform 64-bit floating point operations.

(133) In at least one embodiment, multiple instances of GPGPU 1530 can be configured to operate as a compute cluster. In at least one embodiment, communication used by compute clusters 1536A-1536H for synchronization and data exchange varies across embodiments. In at least one embodiment, multiple instances of GPGPU 1530 communicate over host interface 1532. In at least one embodiment, GPGPU 1530 includes an I/O hub 1539 that couples GPGPU 1530 with a GPU link 1540 that enables a direct connection to other instances of GPGPU 1530. In at least one embodiment, GPU link 1540 is coupled to a dedicated GPU-to-GPU bridge that enables communication and synchronization between multiple instances of GPGPU 1530. In at least one embodiment, GPU link 1540 couples with a high-speed interconnect to transmit and receive data to other GPGPUs or parallel processors. In at least one embodiment, multiple instances of GPGPU 1530 are located in separate data processing systems and communicate via a network device that is accessible via host interface 1532. In at least one embodiment GPU link 1540 can be configured to enable a connection to a host processor in addition to or as an alternative to host interface 1532.

(134) In at least one embodiment, GPGPU 1530 can be configured to train neural networks. In at least one embodiment, GPGPU 1530 can be used within an inferencing platform. In at least one embodiment, in which GPGPU 1530 is used for inferencing, GPGPU 1530 may include fewer compute clusters 1536A-1536H relative to when GPGPU 1530 is used for training a neural network. In at least one embodiment, memory technology associated with memory 1544A-1544B may differ between inferencing and training configurations, with higher bandwidth memory technologies devoted to training configurations. In at least one embodiment, an inferencing

configuration of GPGPU **1530** can support inferencing specific instructions. For example, in at least one embodiment, an inferencing configuration can provide support for one or more 8-bit integer dot product instructions, which may be used during inferencing operations for deployed neural networks.

(135) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in GPGPU **1530** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(136) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(137) FIG. **16** is a block diagram illustrating a computing system **1600** according to at least one embodiment. In at least one embodiment, computing system **1600** includes a processing subsystem **1601** having one or more processor(s) **1602** and a system memory **1604** communicating via an interconnection path that may include a memory hub **1605**. In at least one embodiment, memory hub **1605** may be a separate component within a chipset component or may be integrated within one or more processor(s) **1602**. In at least one embodiment, memory hub **1605** couples with an I/O subsystem **1611** via a communication link **1606**. In at least one embodiment, I/O subsystem **1611** includes an I/O hub **1607** that can enable computing system **1600** to receive input from one or more input device(s) **1608**. In at least one embodiment, I/O hub **1607** can enable a display controller, which may be included in one or more processor(s) **1602**, to provide outputs to one or more display device(s) **1610A**. In at least one embodiment, one or more display device(s) **1610A** coupled with I/O hub **1607** can include a local, internal, or embedded display device.

(138) In at least one embodiment, processing subsystem **1601** includes one or more parallel processor(s) **1612** coupled to memory hub **1605** via a bus or other communication link **1613**. In at least one embodiment, communication link **1613** may use one of any number of standards based communication link technologies or protocols, such as, but not limited to PCI Express, or may be a vendor-specific communications interface or communications fabric. In at least one embodiment, one or more parallel processor(s) **1612** form a computationally focused parallel or vector processing system that can include a large number of processing cores and/or processing clusters, such as a many-integrated core (MIC) processor. In at least one embodiment, some or all of parallel processor(s) **1612** form a graphics processing subsystem that can output pixels to one of one or more display device(s) **1610A** coupled via I/O Hub **1607**. In at least one embodiment, parallel processor(s) **1612** can also include a display controller and display interface (not shown) to enable a direct connection to one or more display device(s) **1610B**. In at least one embodiment, parallel processor(s) **1612** include one or more cores, such as graphics cores **1500** discussed herein.

(139) In at least one embodiment, a system storage unit **1614** can connect to I/O hub **1607** to provide a storage mechanism for computing system **1600**. In at least one embodiment, an I/O switch **1616** can be used to provide an interface mechanism to enable connections between I/O hub **1607** and other components, such as a network adapter **1618** and/or a wireless network adapter **1619** that may be integrated into platform, and various other devices that can be added via one or more add-in device(s) **1620**. In at least one embodiment, network adapter **1618** can be an Ethernet adapter or another wired network adapter. In at least one embodiment, wireless network adapter **1619** can include one or more of a Wi-Fi, Bluetooth, near field communication (NFC), or other network device that includes one or more wireless radios.

(140) In at least one embodiment, computing system **1600** can include other components not explicitly shown, including USB or other port connections, optical storage drives, video capture devices, and like, may also be connected to I/O hub **1607**. In at least one embodiment, communication paths interconnecting various components in FIG. **16** may be implemented using any suitable protocols, such as PCI (Peripheral Component Interconnect) based protocols (e.g.,

PCI-Express), or other bus or point-to-point communication interfaces and/or protocol(s), such as NV-Link high-speed interconnect, or interconnect protocols.

(141) In at least one embodiment, parallel processor(s) **1612** incorporate circuitry optimized for graphics and video processing, including, for example, video output circuitry, and constitutes a graphics processing unit (GPU), e.g., parallel processor(s) **1612** includes graphics core **1500**. In at least one embodiment, parallel processor(s) **1612** incorporate circuitry optimized for general purpose processing. In at least one embodiment, components of computing system **1600** may be integrated with one or more other system elements on a single integrated circuit. For example, in at least one embodiment, parallel processor(s) **1612**, memory hub **1605**, processor(s) **1602**, and I/O hub **1607** can be integrated into a system on chip (SoC) integrated circuit. In at least one embodiment, components of computing system **1600** can be integrated into a single package to form a system in package (SIP) configuration. In at least one embodiment, at least a portion of components of computing system **1600** can be integrated into a multi-chip module (MCM), which can be interconnected with other multi-chip modules into a modular computing system.

(142) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in computing system **1600** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(143) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

Processors

(144) FIG. **17A** illustrates a parallel processor **1700** according to at least one embodiment. In at least one embodiment, various components of parallel processor **1700** may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (ASICs), or field programmable gate arrays (FPGA). In at least one embodiment, illustrated parallel processor **1700** is a variant of one or more parallel processor(s) **1612** shown in FIG. **16** according to an exemplary embodiment. In at least one embodiment, a parallel processor **1700** includes one or more graphics cores **1500**.

(145) In at least one embodiment, parallel processor **1700** includes a parallel processing unit **1702**. In at least one embodiment, parallel processing unit **1702** includes an I/O unit **1704** that enables communication with other devices, including other instances of parallel processing unit **1702**. In at least one embodiment, I/O unit **1704** may be directly connected to other devices. In at least one embodiment, I/O unit **1704** connects with other devices via use of a hub or switch interface, such as a memory hub **1705**. In at least one embodiment, connections between memory hub **1705** and I/O unit **1704** form a communication link **1713**. In at least one embodiment, I/O unit **1704** connects with a host interface **1706** and a memory crossbar **1716**, where host interface **1706** receives commands directed to performing processing operations and memory crossbar **1716** receives commands directed to performing memory operations.

(146) In at least one embodiment, when host interface **1706** receives a command buffer via I/O unit **1704**, host interface **1706** can direct work operations to perform those commands to a front end **1708**. In at least one embodiment, front end **1708** couples with a scheduler **1710** (which may be referred to as a sequencer), which is configured to distribute commands or other work items to a processing cluster array **1712**. In at least one embodiment, scheduler **1710** ensures that processing cluster array **1712** is properly configured and in a valid state before tasks are distributed to a cluster of processing cluster array **1712**. In at least one embodiment, scheduler **1710** is implemented via firmware logic executing on a microcontroller. In at least one embodiment, microcontroller implemented scheduler **1710** is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, enabling rapid preemption and context switching of threads executing on processing array **1712**. In at least one embodiment, host software can prove

workloads for scheduling on processing cluster array **1712** via one of multiple graphics processing paths. In at least one embodiment, workloads can then be automatically distributed across processing array cluster **1712** by scheduler **1710** logic within a microcontroller including scheduler **1710**.

(147) In at least one embodiment, processing cluster array **1712** can include up to “N” processing clusters (e.g., cluster **1714A**, cluster **1714B**, through cluster **1714N**), where “N” represents a positive integer (which may be a different integer “N” than used in other figures). In at least one embodiment, each cluster **1714A-1714N** of processing cluster array **1712** can execute a large number of concurrent threads. In at least one embodiment, scheduler **1710** can allocate work to clusters **1714A-1714N** of processing cluster array **1712** using various scheduling and/or work distribution algorithms, which may vary depending on workload arising for each type of program or computation. In at least one embodiment, scheduling can be handled dynamically by scheduler **1710**, or can be assisted in part by compiler logic during compilation of program logic configured for execution by processing cluster array **1712**. In at least one embodiment, different clusters **1714A-1714N** of processing cluster array **1712** can be allocated for processing different types of programs or for performing different types of computations.

(148) In at least one embodiment, processing cluster array **1712** can be configured to perform various types of parallel processing operations. In at least one embodiment, processing cluster array **1712** is configured to perform general-purpose parallel compute operations. For example, in at least one embodiment, processing cluster array **1712** can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

(149) In at least one embodiment, processing cluster array **1712** is configured to perform parallel graphics processing operations. In at least one embodiment, processing cluster array **1712** can include additional logic to support execution of such graphics processing operations, including but not limited to, texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. In at least one embodiment, processing cluster array **1712** can be configured to execute graphics processing related shader programs such as, but not limited to, vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. In at least one embodiment, parallel processing unit **1702** can transfer data from system memory via I/O unit **1704** for processing. In at least one embodiment, during processing, transferred data can be stored to on-chip memory (e.g., parallel processor memory **1722**) during processing, then written back to system memory.

(150) In at least one embodiment, when parallel processing unit **1702** is used to perform graphics processing, scheduler **1710** can be configured to divide a processing workload into approximately equal sized tasks, to better enable distribution of graphics processing operations to multiple clusters **1714A-1714N** of processing cluster array **1712**. In at least one embodiment, portions of processing cluster array **1712** can be configured to perform different types of processing. For example, in at least one embodiment, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. In at least one embodiment, intermediate data produced by one or more of clusters **1714A-1714N** may be stored in buffers to allow intermediate data to be transmitted between clusters **1714A-1714N** for further processing.

(151) In at least one embodiment, processing cluster array **1712** can receive processing tasks to be executed via scheduler **1710**, which receives commands defining processing tasks from front end **1708**. In at least one embodiment, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and commands defining how data is to be processed (e.g., what program is to be executed). In at least one embodiment, scheduler **1710** may be configured to fetch indices corresponding to tasks or may

receive indices from front end **1708**. In at least one embodiment, front end **1708** can be configured to ensure processing cluster array **1712** is configured to a valid state before a workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

(152) In at least one embodiment, each of one or more instances of parallel processing unit **1702** can couple with a parallel processor memory **1722**. In at least one embodiment, parallel processor memory **1722** can be accessed via memory crossbar **1716**, which can receive memory requests from processing cluster array **1712** as well as I/O unit **1704**. In at least one embodiment, memory crossbar **1716** can access parallel processor memory **1722** via a memory interface **1718**. In at least one embodiment, memory interface **1718** can include multiple partition units (e.g., partition unit **1720A**, partition unit **1720B**, through partition unit **1720N**) that can each couple to a portion (e.g., memory unit) of parallel processor memory **1722**. In at least one embodiment, a number of partition units **1720A-1720N** is configured to be equal to a number of memory units, such that a first partition unit **1720A** has a corresponding first memory unit **1724A**, a second partition unit **1720B** has a corresponding memory unit **1724B**, and an N-th partition unit **1720N** has a corresponding N-th memory unit **1724N**. In at least one embodiment, a number of partition units **1720A-1720N** may not be equal to a number of memory units.

(153) In at least one embodiment, memory units **1724A-1724N** can include various types of memory devices, including dynamic random access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory. In at least one embodiment, memory units **1724A-1724N** may also include 3D stacked memory, including but not limited to high bandwidth memory (HBM), HBM2e, or HBM3. In at least one embodiment, render targets, such as frame buffers or texture maps may be stored across memory units **1724A-1724N**, allowing partition units **1720A-1720N** to write portions of each render target in parallel to efficiently use available bandwidth of parallel processor memory **1722**. In at least one embodiment, a local instance of parallel processor memory **1722** may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

(154) In at least one embodiment, any one of clusters **1714A-1714N** of processing cluster array **1712** can process data that will be written to any of memory units **1724A-1724N** within parallel processor memory **1722**. In at least one embodiment, memory crossbar **1716** can be configured to transfer an output of each cluster **1714A-1714N** to any partition unit **1720A-1720N** or to another cluster **1714A-1714N**, which can perform additional processing operations on an output. In at least one embodiment, each cluster **1714A-1714N** can communicate with memory interface **1718** through memory crossbar **1716** to read from or write to various external memory devices. In at least one embodiment, memory crossbar **1716** has a connection to memory interface **1718** to communicate with I/O unit **1704**, as well as a connection to a local instance of parallel processor memory **1722**, enabling processing units within different processing clusters **1714A-1714N** to communicate with system memory or other memory that is not local to parallel processing unit **1702**. In at least one embodiment, memory crossbar **1716** can use virtual channels to separate traffic streams between clusters **1714A-1714N** and partition units **1720A-1720N**.

(155) In at least one embodiment, multiple instances of parallel processing unit **1702** can be provided on a single add-in card, or multiple add-in cards can be interconnected. In at least one embodiment, different instances of parallel processing unit **1702** can be configured to interoperate even if different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. For example, in at least one embodiment, some instances of parallel processing unit **1702** can include higher precision floating point units relative to other instances. In at least one embodiment, systems incorporating one or more instances of parallel processing unit **1702** or parallel processor **1700** can be implemented in a variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems.

(156) FIG. 17B is a block diagram of a partition unit **1720** according to at least one embodiment. In at least one embodiment, partition unit **1720** is an instance of one of partition units **1720A-1720N** of FIG. 17A. In at least one embodiment, partition unit **1720** includes an L2 cache **1721**, a frame buffer interface **1725**, and a ROP **1726** (raster operations unit). In at least one embodiment, L2 cache **1721** is a read/write cache that is configured to perform load and store operations received from memory crossbar **1716** and ROP **1726**. In at least one embodiment, read misses and urgent write-back requests are output by L2 cache **1721** to frame buffer interface **1725** for processing. In at least one embodiment, updates can also be sent to a frame buffer via frame buffer interface **1725** for processing. In at least one embodiment, frame buffer interface **1725** interfaces with one of memory units in parallel processor memory, such as memory units **1724A-1724N** of FIG. 17A (e.g., within parallel processor memory **1722**).

(157) In at least one embodiment, ROP **1726** is a processing unit that performs raster operations such as stencil, z test, blending, etc. In at least one embodiment, ROP **1726** then outputs processed graphics data that is stored in graphics memory. In at least one embodiment, ROP **1726** includes compression logic to compress depth or color data that is written to memory and decompress depth or color data that is read from memory. In at least one embodiment, compression logic can be lossless compression logic that makes use of one or more of multiple compression algorithms. In at least one embodiment, a type of compression that is performed by ROP **1726** can vary based on statistical characteristics of data to be compressed. For example, in at least one embodiment, delta color compression is performed on depth and color data on a per-tile basis.

(158) In at least one embodiment, ROP **1726** is included within each processing cluster (e.g., cluster **1714A-1714N** of FIG. 17A) instead of within partition unit **1720**. In at least one embodiment, read and write requests for pixel data are transmitted over memory crossbar **1716** instead of pixel fragment data. In at least one embodiment, processed graphics data may be displayed on a display device, such as one of one or more display device(s) **1610** of FIG. 16, routed for further processing by processor(s) **1602**, or routed for further processing by one of processing entities within parallel processor **1700** of FIG. 17A.

(159) FIG. 18 illustrates a multi-GPU computing system **1800**, according to at least one embodiment. In at least one embodiment, multi-GPU computing system **1800** can include a processor **1802** coupled to multiple general purpose graphics processing units (GPGPUs) **1806A-D** via a host interface switch **1804**. In at least one embodiment, host interface switch **1804** is a PCI express switch device that couples processor **1802** to a PCI express bus over which processor **1802** can communicate with GPGPUs **1806A-D**. In at least one embodiment, GPGPUs **1806A-D** can interconnect via a set of high-speed point-to-point GPU-to-GPU links **1816**. In at least one embodiment, GPU-to-GPU links **1816** connect to each of GPGPUs **1806A-D** via a dedicated GPU link. In at least one embodiment, P2P GPU links **1816** enable direct communication between each of GPGPUs **1806A-D** without requiring communication over host interface bus **1804** to which processor **1802** is connected. In at least one embodiment, with GPU-to-GPU traffic directed to P2P GPU links **1816**, host interface bus **1804** remains available for system memory access or to communicate with other instances of multi-GPU computing system **1800**, for example, via one or more network devices. While in at least one embodiment GPGPUs **1806A-D** connect to processor **1802** via host interface switch **1804**, in at least one embodiment processor **1802** includes direct support for P2P GPU links **1816** and can connect directly to GPGPUs **1806A-D**.

(160) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in multi-GPU computing system **1800** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein. In at least one embodiment, multi-GPU computing system **1800** includes one or more graphics cores **1500**.

(161) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(162) FIG. **19** is a block diagram of a graphics processor **1900**, according to at least one embodiment. In at least one embodiment, graphics processor **1900** includes a ring interconnect **1902**, a pipeline front-end **1904**, a media engine **1937**, and graphics cores **1980A-1980N**. In at least one embodiment, ring interconnect **1902** couples graphics processor **1900** to other processing units, including other graphics processors or one or more general-purpose processor cores. In at least one embodiment, graphics processor **1900** is one of many processors integrated within a multi-core processing system. In at least one embodiment, graphics processor **1900** includes graphics core **1500**.

(163) In at least one embodiment, graphics processor **1900** receives batches of commands via ring interconnect **1902**. In at least one embodiment, incoming commands are interpreted by a command streamer **1903** in pipeline front-end **1904**. In at least one embodiment, graphics processor **1900** includes scalable execution logic to perform 3D geometry processing and media processing via graphics core(s) **1980A-1980N**. In at least one embodiment, for 3D geometry processing commands, command streamer **1903** supplies commands to geometry pipeline **1936**. In at least one embodiment, for at least some media processing commands, command streamer **1903** supplies commands to a video front end **1934**, which couples with media engine **1937**. In at least one embodiment, media engine **1937** includes a Video Quality Engine (VQE) **1930** for video and image post-processing and a multi-format encode/decode (MFX) **1933** engine to provide hardware-accelerated media data encoding and decoding. In at least one embodiment, geometry pipeline **1936** and media engine **1937** each generate execution threads for thread execution resources provided by at least one graphics core **1980**.

(164) In at least one embodiment, graphics processor **1900** includes scalable thread execution resources featuring graphics cores **1980A-1980N** (which can be modular and are sometimes referred to as core slices), each having multiple sub-cores **1950A-50N**, **1960A-1960N** (sometimes referred to as core sub-slices). In at least one embodiment, graphics processor **1900** can have any number of graphics cores **1980A**. In at least one embodiment, graphics processor **1900** includes a graphics core **1980A** having at least a first sub-core **1950A** and a second sub-core **1960A**. In at least one embodiment, graphics processor **1900** is a low power processor with a single sub-core (e.g., **1950A**). In at least one embodiment, graphics processor **1900** includes multiple graphics cores **1980A-1980N**, each including a set of first sub-cores **1950A-1950N** and a set of second sub-cores **1960A-1960N**. In at least one embodiment, each sub-core in first sub-cores **1950A-1950N** includes at least a first set of execution units **1952A-1952N** and media/texture samplers **1954A-1954N**. In at least one embodiment, each sub-core in second sub-cores **1960A-1960N** includes at least a second set of execution units **1962A-1962N** and samplers **1964A-1964N**. In at least one embodiment, each sub-core **1950A-1950N**, **1960A-1960N** shares a set of shared resources **1970A-1970N**. In at least one embodiment, shared resources include shared cache memory and pixel operation logic. In at least one embodiment, graphics processor **1900** includes load/store units in pipeline front-end **1904**.

(165) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment, inference and/or training logic **715** may be used in graphics processor **1900** for inferencing or predicting operations based, at least in part, on weight parameters calculated using neural network training operations, neural network functions and/or architectures, or neural network use cases described herein.

(166) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(167) FIG. **20** is a block diagram of a processing system, according to at least one embodiment. In at least one embodiment, system **2000** includes one or more processors **2002** and one or more graphics processors **2008**, and may be a single processor desktop system, a multiprocessor

workstation system, or a server system having a large number of processors **2002** or processor cores **2007**. In at least one embodiment, system **2000** is a processing platform incorporated within a system-on-a-chip (SoC) integrated circuit for use in mobile, handheld, or embedded devices. In at least one embodiment, one or more graphics processors **2008** include one or more graphics cores **1500**.

(168) In at least one embodiment, system **2000** can include, or be incorporated within a server-based gaming platform, a game console, including a game and media console, a mobile gaming console, a handheld game console, or an online game console. In at least one embodiment, system **2000** is a mobile phone, a smart phone, a tablet computing device or a mobile Internet device. In at least one embodiment, processing system **2000** can also include, couple with, or be integrated within a wearable device, such as a smart watch wearable device, a smart eyewear device, an augmented reality device, or a virtual reality device. In at least one embodiment, processing system **2000** is a television or set top box device having one or more processors **2002** and a graphical interface generated by one or more graphics processors **2008**.

(169) In at least one embodiment, one or more processors **2002** each include one or more processor cores **2007** to process instructions which, when executed, perform operations for system and user software. In at least one embodiment, each of one or more processor cores **2007** is configured to process a specific instruction sequence **2009**. In at least one embodiment, instruction sequence **2009** may facilitate Complex Instruction Set Computing (CISC), Reduced Instruction Set Computing (RISC), or computing via a Very Long Instruction Word (VLIW). In at least one embodiment, processor cores **2007** may each process a different instruction sequence **2009**, which may include instructions to facilitate emulation of other instruction sequences. In at least one embodiment, processor core **2007** may also include other processing devices, such a Digital Signal Processor (DSP).

(170) In at least one embodiment, processor **2002** includes a cache memory **2004**. In at least one embodiment, processor **2002** can have a single internal cache or multiple levels of internal cache. In at least one embodiment, cache memory is shared among various components of processor **2002**. In at least one embodiment, processor **2002** also uses an external cache (e.g., a Level-3 (L3) cache or Last Level Cache (LLC)) (not shown), which may be shared among processor cores **2007** using known cache coherency techniques. In at least one embodiment, a register file **2006** is additionally included in processor **2002**, which may include different types of registers for storing different types of data (e.g., integer registers, floating point registers, status registers, and an instruction pointer register). In at least one embodiment, register file **2006** may include general-purpose registers or other registers.

(171) In at least one embodiment, one or more processor(s) **2002** are coupled with one or more interface bus(es) **2010** to transmit communication signals such as address, data, or control signals between processor **2002** and other components in system **2000**. In at least one embodiment, interface bus **2010** can be a processor bus, such as a version of a Direct Media Interface (DMI) bus. In at least one embodiment, interface bus **2010** is not limited to a DMI bus, and may include one or more Peripheral Component Interconnect buses (e.g., PCI, PCI Express), memory busses, or other types of interface busses. In at least one embodiment processor(s) **2002** include an integrated memory controller **2016** and a platform controller hub **2030**. In at least one embodiment, memory controller **2016** facilitates communication between a memory device and other components of system **2000**, while platform controller hub (PCH) **2030** provides connections to I/O devices via a local I/O bus.

(172) In at least one embodiment, a memory device **2020** can be a dynamic random access memory (DRAM) device, a static random access memory (SRAM) device, flash memory device, phase-change memory device, or some other memory device having suitable performance to serve as process memory. In at least one embodiment, memory device **2020** can operate as system memory for system **2000**, to store data **2022** and instructions **2021** for use when one or more processors

2002 executes an application or process. In at least one embodiment, memory controller **2016** also couples with an optional external graphics processor **2012**, which may communicate with one or more graphics processors **2008** in processors **2002** to perform graphics and media operations. In at least one embodiment, a display device **2011** can connect to processor(s) **2002**. In at least one embodiment, display device **2011** can include one or more of an internal display device, as in a mobile electronic device or a laptop device, or an external display device attached via a display interface (e.g., DisplayPort, etc.). In at least one embodiment, display device **2011** can include a head mounted display (HMD) such as a stereoscopic display device for use in virtual reality (VR) applications or augmented reality (AR) applications.

(173) In at least one embodiment, platform controller hub **2030** enables peripherals to connect to memory device **2020** and processor **2002** via a high-speed I/O bus. In at least one embodiment, I/O peripherals include, but are not limited to, an audio controller **2046**, a network controller **2034**, a firmware interface **2028**, a wireless transceiver **2026**, touch sensors **2025**, a data storage device **2024** (e.g., hard disk drive, flash memory, etc.). In at least one embodiment, data storage device **2024** can connect via a storage interface (e.g., SATA) or via a peripheral bus, such as a Peripheral Component Interconnect bus (e.g., PCI, PCI Express). In at least one embodiment, touch sensors **2025** can include touch screen sensors, pressure sensors, or fingerprint sensors. In at least one embodiment, wireless transceiver **2026** can be a Wi-Fi transceiver, a Bluetooth transceiver, or a mobile network transceiver such as a 3G, 4G, or Long Term Evolution (LTE) transceiver. In at least one embodiment, firmware interface **2028** enables communication with system firmware, and can be, for example, a unified extensible firmware interface (UEFI). In at least one embodiment, network controller **2034** can enable a network connection to a wired network. In at least one embodiment, a high-performance network controller (not shown) couples with interface bus **2010**. In at least one embodiment, audio controller **2046** is a multi-channel high definition audio controller. In at least one embodiment, system **2000** includes an optional legacy I/O controller **2040** for coupling legacy (e.g., Personal System 2 (PS/2)) devices to system **2000**. In at least one embodiment, platform controller hub **2030** can also connect to one or more Universal Serial Bus (USB) controllers **2042** connect input devices, such as keyboard and mouse **2043** combinations, a camera **2044**, or other USB input devices.

(174) In at least one embodiment, an instance of memory controller **2016** and platform controller hub **2030** may be integrated into a discreet external graphics processor, such as external graphics processor **2012**. In at least one embodiment, platform controller hub **2030** and/or memory controller **2016** may be external to one or more processor(s) **2002**. For example, in at least one embodiment, system **2000** can include an external memory controller **2016** and platform controller hub **2030**, which may be configured as a memory controller hub and peripheral controller hub within a system chipset that is in communication with processor(s) **2002**.

(175) Inference and/or training logic **715** are used to perform inferencing and/or training operations associated with one or more embodiments. In at least one embodiment portions or all of inference and/or training logic **715** may be incorporated into graphics processor **2008**. For example, in at least one embodiment, training and/or inferencing techniques described herein may use one or more of ALUs embodied in a 3D pipeline. In at least one embodiment, weight parameters may be stored in on-chip or off-chip memory and/or registers (shown or not shown) that configure ALUs of graphics processor **2008** to perform one or more machine learning algorithms, neural network architectures, use cases, or training techniques described herein.

(176) Embodiments presented herein can allow for session keys to be wrapped using a sequence of wrapping keys to provide a secure tunnel for transmission of the wrapped session keys

(177) At least one embodiment of the disclosure can be described in view of the following clauses:
1. A method, comprising: generating a session key using a software agent executing on a secure processor in a computing device; encrypting the session key using a wrapping key generated using the secure processor; transmitting the session key, encrypted using the wrapping key, over a system

bus in the computing device; decrypting, at a secure key unwrap engine implemented in hardware of the computing device and receiving the session key, the session key using the wrapping key; and delivering the session key, in plaintext, to an inline crypto engine executing on the computing device, wherein the inline crypto-engine is allowed to use the session key to encrypt session data to be transmitted to a second device. 2. The method of clause 1, wherein the wrapping key is used to encrypt a group of session keys for a plurality of sessions. 3. The method of clause 1, further comprising: transmitting, over the system bus, one or more new wrapping keys to be used by the crypto-engine to encrypt future session keys. 4. The method of clause 3, wherein a key manager executing on the secure processor determines the times at which to transmit the new wrapping keys and generates the new wrapping keys to be transmitted. 5. The method of clause 1, further comprising: receiving an initial wrapping key agreed to according to a communication protocol; encrypting the wrapping key using the initial wrapping key; and transmitting the wrapping key, encrypted using the initial wrapping key, over the system bus to be decrypted by the secure key unwrap engine, wherein the initial wrapping key is to be received without encryption during a boot process of the computing device. 6. The method of clause 1, wherein the software agent is one of a plurality of protocol-specific software agents executing within a central root of trust on the computing device to generate session keys to be used for encryption by one of a plurality of protocol-specific inline crypto-engines. 7. The method of clause 1, further comprising: analyzing, using the key unwrap engine, metadata associated with the session key to determine information about the session key, the information indicating at least one of a location to route the session key, an identifier for the session key, or a type of the session key. 8. The method of clause 1, further comprising: analyzing, using the key unwrap engine, a status bit for a received key to determine whether a type of the received key is a session key type or a wrapping key type. 9. The method of clause 1, wherein the computing device is a network interface card (NIC) or a graphics card including one or more graphics processing units. 10. A system, comprising: a secure processor to generate a session key using a software agent executing on the secure processor, and further to encrypt the session key using a wrapping key generated using the secure processor; a system bus to transmit the session key wrapped by the wrapping key; a key unwrap engine, implemented in hardware on the system, to receive the session key over the system bus and decrypt the session key using the wrapping key; and a crypto engine receive the decrypted session key, wherein the crypto engine is allowed to encrypt session data using the decrypted session key to be transmitted to a second device. 11. The system of clause 10, wherein the wrapping key is used to encrypt a group of session keys for a plurality of sessions. 12. The system of clause 10, wherein the secure process is further to transmit, at a series of times and over the system bus, new wrapping keys to be used by the crypto-engine to encrypt future session keys. 13. The system of clause 10, wherein the software agent is one of a plurality of protocol-specific software agents executing within a central root of trust on the system to generate session keys to be used for encryption by one of a plurality of protocol-specific inline crypto-engines executing on the system. 14. The system of clause 10, wherein the secure processor is further to receive an initial wrapping key and encrypt the wrapping key using the initial wrapping key, the secure processor further to transmit the wrapping key, encrypted using the initial wrapping key, over the system bus to be decrypted by the secure key unwrap engine and transmitted to the crypto-engine, wherein the initial wrapping key is to be received without encryption during a boot process of the computing device. 15. The system of clause 10, wherein the key unwrap engine is further to analyze metadata associated with the session key to determine information about the session key, the information indicating at least one of a location to route the session key, an identifier for the session key, or a type of the session key. 16. A key tunnel including one or more processors for receiving session keys transmitted over a system bus of a computing device, the session keys encrypted within a root of trust on the computing device using a wrapping key generated within the root of trust, the one or more processors of the key tunnel to unwrap the session keys using the wrapping key and transmit the session keys to one

or more protocol-specific crypto engines for encrypting session data to be transmitted to one or more remote devices. 17. The key tunnel of clause 16, wherein the one or more processors are further to analyze metadata associated with the session key to determine where to route the session key. 18. The key tunnel of clause 16, wherein the one or more processors are further to analyze a status bit for a received key to determine whether a type of the received key is received key is a session key type or a wrapping key type. 19. The key tunnel of clause 16, wherein the root of trust is to use an initial shared secret to encrypt the wrapping key to be received by the key tunnel, wherein the initial shared secret is one of a plaintext key, an embedded key, a provisioned key shared by two or more fuse blocks, a secret received over a physical private bus, or a secret generated using a specified cryptographic scheme. 20. The key tunnel of clause 16, wherein the key tunnel is comprised in at least one of: a system for performing simulation operations; a system for performing simulation operations to test or validate autonomous machine applications; a system for rendering graphical output; a system for performing deep learning operations; a system implemented using an edge device; a system for generating or presenting virtual reality (VR) content; a system for generating or presenting augmented reality (AR) content; a system for generating or presenting mixed reality (MR) content; a system incorporating one or more Virtual Machines (VMs); a system implemented at least partially in a data center; a system for performing hardware testing using simulation; a system for synthetic data generation; a collaborative content creation platform for 3D assets; or a system implemented at least partially using cloud computing resources.

(178) In at least one embodiment, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit or chip. In at least one embodiment, multi-chip modules may be used with increased connectivity which simulate on-chip operation, and make substantial improvements over utilizing a conventional central processing unit (“CPU”) and bus implementation. In at least one embodiment, various modules may also be situated separately or in various combinations of semiconductor platforms per desires of user.

(179) In at least one embodiment, referring back to FIG. **10**, computer programs in form of machine-readable executable code or computer control logic algorithms are stored in main memory **1004** and/or secondary storage. Computer programs, if executed by one or more processors, enable system **1000** to perform various functions in accordance with at least one embodiment. In at least one embodiment, memory **1004**, storage, and/or any other storage are possible examples of computer-readable media. In at least one embodiment, secondary storage may refer to any suitable storage device or system such as a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (“DVD”) drive, recording device, universal serial bus (“USB”) flash memory, etc. In at least one embodiment, architecture and/or functionality of various previous figures are implemented in context of CPU **1002**, parallel processing system **1012**, an integrated circuit capable of at least a portion of capabilities of both CPU **1002**, parallel processing system **1012**, a chipset (e.g., a group of integrated circuits designed to work and sold as a unit for performing related functions, etc.), and/or any suitable combination of integrated circuit(s).

(180) In at least one embodiment, architecture and/or functionality of various previous figures are implemented in context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and more. In at least one embodiment, computer system **1000** may take form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (“PDA”), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(181) In at least one embodiment, parallel processing system **1012** includes, without limitation, a plurality of parallel processing units (“PPUs”) **1014** and associated memories **1016**. In at least one

embodiment, PPU**s** **1014** are connected to a host processor or other peripheral devices via an interconnect **1018** and a switch **1020** or multiplexer. In at least one embodiment, parallel processing system **1012** distributes computational tasks across PPU**s** **1014** which can be parallelizable—for example, as part of distribution of computational tasks across multiple graphics processing unit (“GPU”) thread blocks. In at least one embodiment, memory is shared and accessible (e.g., for read and/or write access) across some or all of PPU**s** **1014**, although such shared memory may incur performance penalties relative to use of local memory and registers resident to a PPU **1014**. In at least one embodiment, operation of PPU**s** **1014** is synchronized through use of a command such as `_syncthreads( )`, wherein all threads in a block (e.g., executed across multiple PPU**s** **1014**) to reach a certain point of execution of code before proceeding.

(182) In at least one embodiment, one or more techniques described herein utilize a oneAPI programming model. In at least one embodiment, a oneAPI programming model refers to a programming model for interacting with various compute accelerator architectures. In at least one embodiment, oneAPI refers to an application programming interface (API) designed to interact with various compute accelerator architectures. In at least one embodiment, a oneAPI programming model utilizes a DPC++ programming language. In at least one embodiment, a DPC++ programming language refers to a high-level language for data parallel programming productivity. In at least one embodiment, a DPC++ programming language is based at least in part on C and/or C++ programming languages. In at least one embodiment, a oneAPI programming model is a programming model such as those developed by Intel Corporation of Santa Clara, CA.

(183) In at least one embodiment, oneAPI and/or oneAPI programming model is utilized to interact with various accelerator, GPU, processor, and/or variations thereof, architectures. In at least one embodiment, oneAPI includes a set of libraries that implement various functionalities. In at least one embodiment, oneAPI includes at least a oneAPI DPC++ library, a oneAPI math kernel library, a oneAPI data analytics library, a oneAPI deep neural network library, a oneAPI collective communications library, a oneAPI threading building blocks library, a oneAPI video processing library, and/or variations thereof.

(184) In at least one embodiment, a oneAPI DPC++ library, also referred to as oneDPL, is a library that implements algorithms and functions to accelerate DPC++ kernel programming. In at least one embodiment, oneDPL implements one or more standard template library (STL) functions. In at least one embodiment, oneDPL implements one or more parallel STL functions. In at least one embodiment, oneDPL provides a set of library classes and functions such as parallel algorithms, iterators, function object classes, range-based API, and/or variations thereof. In at least one embodiment, oneDPL implements one or more classes and/or functions of a C++ standard library. In at least one embodiment, oneDPL implements one or more random number generator functions.

(185) In at least one embodiment, a oneAPI math kernel library, also referred to as oneMKL, is a library that implements various optimized and parallelized routines for various mathematical functions and/or operations. In at least one embodiment, oneMKL implements one or more basic linear algebra subprograms (BLAS) and/or linear algebra package (LAPACK) dense linear algebra routines. In at least one embodiment, oneMKL implements one or more sparse BLAS linear algebra routines. In at least one embodiment, oneMKL implements one or more random number generators (RNGs). In at least one embodiment, oneMKL implements one or more vector mathematics (VM) routines for mathematical operations on vectors. In at least one embodiment, oneMKL implements one or more Fast Fourier Transform (FFT) functions.

(186) In at least one embodiment, a oneAPI data analytics library, also referred to as oneDAL, is a library that implements various data analysis applications and distributed computations. In at least one embodiment, oneDAL implements various algorithms for preprocessing, transformation, analysis, modeling, validation, and decision making for data analytics, in batch, online, and distributed processing modes of computation. In at least one embodiment, oneDAL implements various C++ and/or Java APIs and various connectors to one or more data sources. In at least one

embodiment, oneDAL implements DPC++ API extensions to a traditional C++ interface and enables GPU usage for various algorithms.

(187) In at least one embodiment, a oneAPI deep neural network library, also referred to as oneDNN, is a library that implements various deep learning functions. In at least one embodiment, oneDNN implements various neural network, machine learning, and deep learning functions, algorithms, and/or variations thereof.

(188) In at least one embodiment, a oneAPI collective communications library, also referred to as oneCCL, is a library that implements various applications for deep learning and machine learning workloads. In at least one embodiment, oneCCL is built upon lower-level communication middleware, such as message passing interface (MPI) and libfabrics. In at least one embodiment, oneCCL enables a set of deep learning specific optimizations, such as prioritization, persistent operations, out of order executions, and/or variations thereof. In at least one embodiment, oneCCL implements various CPU and GPU functions.

(189) In at least one embodiment, a oneAPI threading building blocks library, also referred to as oneTBB, is a library that implements various parallelized processes for various applications. In at least one embodiment, oneTBB is utilized for task-based, shared parallel programming on a host. In at least one embodiment, oneTBB implements generic parallel algorithms. In at least one embodiment, oneTBB implements concurrent containers. In at least one embodiment, oneTBB implements a scalable memory allocator. In at least one embodiment, oneTBB implements a work-stealing task scheduler. In at least one embodiment, oneTBB implements low-level synchronization primitives. In at least one embodiment, oneTBB is compiler-independent and usable on various processors, such as GPUs, PPU, CPUs, and/or variations thereof.

(190) In at least one embodiment, a oneAPI video processing library, also referred to as oneVPL, is a library that is utilized for accelerating video processing in one or more applications. In at least one embodiment, oneVPL implements various video decoding, encoding, and processing functions. In at least one embodiment, oneVPL implements various functions for media pipelines on CPUs, GPUs, and other accelerators. In at least one embodiment, oneVPL implements device discovery and selection in media centric and video analytics workloads. In at least one embodiment, oneVPL implements API primitives for zero-copy buffer sharing.

(191) In at least one embodiment, a oneAPI programming model utilizes a DPC++ programming language. In at least one embodiment, a DPC++ programming language is a programming language that includes, without limitation, functionally similar versions of CUDA mechanisms to define device code and distinguish between device code and host code. In at least one embodiment, a DPC++ programming language may include a subset of functionality of a CUDA programming language. In at least one embodiment, one or more CUDA programming model operations are performed using a oneAPI programming model using a DPC++ programming language.

(192) In at least one embodiment, any application programming interface (API) described herein is compiled into one or more instructions, operations, or any other signal by a compiler, interpreter, or other software tool. In at least one embodiment, compilation comprises generating one or more machine-executable instructions, operations, or other signals from source code. In at least one embodiment, an API compiled into one or more instructions, operations, or other signals, when performed, causes one or more processors such as graphics processors **1410** and **1440**, graphics cores **1500**, parallel processor **1700**, processors **1802**, processors **2002**, or any other logic circuit further described herein to perform one or more computing operations.

(193) It should be noted that, while example embodiments described herein may relate to a CUDA programming model, techniques described herein can be utilized with any suitable programming model, such as HIP, oneAPI, and/or variations thereof.

(194) Other variations are within spirit of present disclosure. Thus, while disclosed techniques are susceptible to various modifications and alternative constructions, certain illustrated embodiments thereof are shown in drawings and have been described above in detail. It should be understood,

however, that there is no intention to limit disclosure to specific form or forms disclosed, but on contrary, intention is to cover all modifications, alternative constructions, and equivalents falling within spirit and scope of disclosure, as defined in appended claims.

(195) Use of terms “a” and “an” and “the” and similar referents in context of describing disclosed embodiments (especially in context of following claims) are to be construed to cover both singular and plural, unless otherwise indicated herein or clearly contradicted by context, and not as a definition of a term. Terms “comprising,” “having,” “including,” and “containing” are to be construed as open-ended terms (meaning “including, but not limited to,”) unless otherwise noted. “Connected,” when unmodified and referring to physical connections, is to be construed as partly or wholly contained within, attached to, or joined together, even if there is something intervening. Recitation of ranges of values herein are merely intended to serve as a shorthand method of referring individually to each separate value falling within range, unless otherwise indicated herein and each separate value is incorporated into specification as if it were individually recited herein. In at least one embodiment, use of term “set” (e.g., “a set of items”) or “subset” unless otherwise noted or contradicted by context, is to be construed as a nonempty collection comprising one or more members. Further, unless otherwise noted or contradicted by context, term “subset” of a corresponding set does not necessarily denote a proper subset of corresponding set, but subset and corresponding set may be equal.

(196) Conjunctive language, such as phrases of form “at least one of A, B, and C,” or “at least one of A, B and C,” unless specifically stated otherwise or otherwise clearly contradicted by context, is otherwise understood with context as used in general to present that an item, term, etc., may be either A or B or C, or any nonempty subset of set of A and B and C. For instance, in illustrative example of a set having three members, conjunctive phrases “at least one of A, B, and C” and “at least one of A, B and C” refer to any of following sets: {A}, {B}, {C}, {A, B}, {A, C}, {B, C}, {A, B, C}. Thus, such conjunctive language is not generally intended to imply that certain embodiments require at least one of A, at least one of B and at least one of C each to be present. In addition, unless otherwise noted or contradicted by context, term “plurality” indicates a state of being plural (e.g., “a plurality of items” indicates multiple items). In at least one embodiment, number of items in a plurality is at least two, but can be more when so indicated either explicitly or by context. Further, unless stated otherwise or otherwise clear from context, phrase “based on” means “based at least in part on” and not “based solely on.”

(197) Operations of processes described herein can be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context. In at least one embodiment, a process such as those processes described herein (or variations and/or combinations thereof) is performed under control of one or more computer systems configured with executable instructions and is implemented as code (e.g., executable instructions, one or more computer programs or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. In at least one embodiment, code is stored on a computer-readable storage medium, for example, in form of a computer program comprising a plurality of instructions executable by one or more processors. In at least one embodiment, a computer-readable storage medium is a non-transitory computer-readable storage medium that excludes transitory signals (e.g., a propagating transient electric or electromagnetic transmission) but includes non-transitory data storage circuitry (e.g., buffers, cache, and queues) within transceivers of transitory signals. In at least one embodiment, code (e.g., executable code or source code) is stored on a set of one or more non-transitory computer-readable storage media having stored thereon executable instructions (or other memory to store executable instructions) that, when executed (i.e., as a result of being executed) by one or more processors of a computer system, cause computer system to perform operations described herein. In at least one embodiment, set of non-transitory computer-readable storage media comprises multiple non-transitory computer-readable storage media and one or more of individual non-transitory storage media of multiple non-

transitory computer-readable storage media lack all of code while multiple non-transitory computer-readable storage media collectively store all of code. In at least one embodiment, executable instructions are executed such that different instructions are executed by different processors—for example, a non-transitory computer-readable storage medium store instructions and a main central processing unit (“CPU”) executes some of instructions while a graphics processing unit (“GPU”) executes other instructions. In at least one embodiment, different components of a computer system have separate processors and different processors execute different subsets of instructions.

(198) In at least one embodiment, an arithmetic logic unit is a set of combinational logic circuitry that takes one or more inputs to produce a result. In at least one embodiment, an arithmetic logic unit is used by a processor to implement mathematical operation such as addition, subtraction, or multiplication. In at least one embodiment, an arithmetic logic unit is used to implement logical operations such as logical AND/OR or XOR. In at least one embodiment, an arithmetic logic unit is stateless, and made from physical switching components such as semiconductor transistors arranged to form logical gates. In at least one embodiment, an arithmetic logic unit may operate internally as a stateful logic circuit with an associated clock. In at least one embodiment, an arithmetic logic unit may be constructed as an asynchronous logic circuit with an internal state not maintained in an associated register set. In at least one embodiment, an arithmetic logic unit is used by a processor to combine operands stored in one or more registers of the processor and produce an output that can be stored by the processor in another register or a memory location.

(199) In at least one embodiment, as a result of processing an instruction retrieved by the processor, the processor presents one or more inputs or operands to an arithmetic logic unit, causing the arithmetic logic unit to produce a result based at least in part on an instruction code provided to inputs of the arithmetic logic unit. In at least one embodiment, the instruction codes provided by the processor to the ALU are based at least in part on the instruction executed by the processor. In at least one embodiment combinational logic in the ALU processes the inputs and produces an output which is placed on a bus within the processor. In at least one embodiment, the processor selects a destination register, memory location, output device, or output storage location on the output bus so that clocking the processor causes the results produced by the ALU to be sent to the desired location.

(200) In the scope of this application, the term arithmetic logic unit, or ALU, is used to refer to any computational logic circuit that processes operands to produce a result. For example, in the present document, the term ALU can refer to a floating point unit, a DSP, a tensor core, a shader core, a coprocessor, or a CPU.

(201) Accordingly, in at least one embodiment, computer systems are configured to implement one or more services that singly or collectively perform operations of processes described herein and such computer systems are configured with applicable hardware and/or software that enable performance of operations. Further, a computer system that implements at least one embodiment of present disclosure is a single device and, in another embodiment, is a distributed computer system comprising multiple devices that operate differently such that distributed computer system performs operations described herein and such that a single device does not perform all operations.

(202) Use of any and all examples, or exemplary language (e.g., “such as”) provided herein, is intended merely to better illuminate embodiments of disclosure and does not pose a limitation on scope of disclosure unless otherwise claimed. No language in specification should be construed as indicating any non-claimed element as essential to practice of disclosure.

(203) All references, including publications, patent applications, and patents, cited herein are hereby incorporated by reference to same extent as if each reference were individually and specifically indicated to be incorporated by reference and were set forth in its entirety herein.

(204) In description and claims, terms “coupled” and “connected,” along with their derivatives, may be used. It should be understood that these terms may be not intended as synonyms for each

other. Rather, in particular examples, “connected” or “coupled” may be used to indicate that two or more elements are in direct or indirect physical or electrical contact with each other. “Coupled” may also mean that two or more elements are not in direct contact with each other, but yet still cooperate or interact with each other.

(205) Unless specifically stated otherwise, it may be appreciated that throughout specification terms such as “processing,” “computing,” “calculating,” “determining,” or like, refer to action and/or processes of a computer or computing system, or similar electronic computing device, that manipulate and/or transform data represented as physical, such as electronic, quantities within computing system's registers and/or memories into other data similarly represented as physical quantities within computing system's memories, registers or other such information storage, transmission or display devices.

(206) In a similar manner, term “processor” may refer to any device or portion of a device that processes electronic data from registers and/or memory and transform that electronic data into other electronic data that may be stored in registers and/or memory. As non-limiting examples, “processor” may be a CPU or a GPU. A “computing platform” may comprise one or more processors. As used herein, “software” processes may include, for example, software and/or hardware entities that perform work over time, such as tasks, threads, and intelligent agents. Also, each process may refer to multiple processes, for carrying out instructions in sequence or in parallel, continuously or intermittently. In at least one embodiment, terms “system” and “method” are used herein interchangeably insofar as system may embody one or more methods and methods may be considered a system.

(207) In present document, references may be made to obtaining, acquiring, receiving, or inputting analog or digital data into a subsystem, computer system, or computer-implemented machine. In at least one embodiment, process of obtaining, acquiring, receiving, or inputting analog and digital data can be accomplished in a variety of ways such as by receiving data as a parameter of a function call or a call to an application programming interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a serial or parallel interface. In at least one embodiment, processes of obtaining, acquiring, receiving, or inputting analog or digital data can be accomplished by transferring data via a computer network from providing entity to acquiring entity. In at least one embodiment, references may also be made to providing, outputting, transmitting, sending, or presenting analog or digital data. In various examples, processes of providing, outputting, transmitting, sending, or presenting analog or digital data can be accomplished by transferring data as an input or output parameter of a function call, a parameter of an application programming interface or interprocess communication mechanism.

(208) Although descriptions herein set forth example implementations of described techniques, other architectures may be used to implement described functionality, and are intended to be within scope of this disclosure. Furthermore, although specific distributions of responsibilities may be defined above for purposes of description, various functions and responsibilities might be distributed and divided in different ways, depending on circumstances.

(209) Furthermore, although subject matter has been described in language specific to structural features and/or methodological acts, it is to be understood that subject matter claimed in appended claims is not necessarily limited to specific features or acts described. Rather, specific features and acts are disclosed as exemplary forms of implementing the claims.

Claims

1. A method, comprising: generating a session key using a software agent executing on a secure processor in a computing device; encrypting the session key using a wrapping key generated using the secure processor; transmitting the session key, encrypted using the wrapping key, over a system

bus to a secure key unwrap engine implemented in a hardware of the computing device; decrypting, at the secure key unwrap engine, the session key using the wrapping key; and delivering the session key, in plaintext, to an inline crypto engine executing on the computing device, wherein the inline crypto engine is allowed to use the session key to encrypt session data to be transmitted to a second device.

2. The method of claim 1, wherein the wrapping key is used to encrypt a group of session keys for a plurality of sessions.

3. The method of claim 1, further comprising: transmitting, over the system bus, one or more new wrapping keys to be used by the crypto engine to encrypt future session keys.

4. The method of claim 3, wherein a key manager executing on the secure processor determines one or more times at which to transmit the new wrapping keys and generates the new wrapping keys to be transmitted.

5. The method of claim 1, further comprising: receiving an initial wrapping key agreed to according to a communication protocol; encrypting the wrapping key using the initial wrapping key; and transmitting the wrapping key, encrypted using the initial wrapping key, over the system bus to be decrypted by the secure key unwrap engine, wherein the initial wrapping key is to be received without encryption during an early execution process of the computing device.

6. The method of claim 1, wherein the software agent is one of a plurality of protocol-specific software agents executing within a central root of trust on the computing device to generate session keys to be used for encryption by one of a plurality of protocol-specific inline crypto engines.

7. The method of claim 1, further comprising: analyzing, using the key unwrap engine, metadata associated with the session key to determine information about the session key, the information indicating at least one of a location to route the session key, an identifier for the session key, or a type of the session key.

8. The method of claim 1, further comprising: analyzing, using the key unwrap engine, a status bit for a received key to determine whether a type of the received key is a session key type or a wrapping key type.

9. The method of claim 1, wherein the computing device is a network interface card (NIC) or a graphics card including one or more graphics processing units.

10. A system, comprising: a secure processor to generate a session key using a software agent executing on the secure processor, and further to encrypt the session key using a wrapping key generated using the secure processor; a system bus to transmit the session key wrapped by the wrapping key within hardware on the system; a key unwrap engine, implemented in the hardware on the system, to receive the session key over the system bus and decrypt the session key using the wrapping key; and a crypto engine to receive the decrypted session key, wherein the crypto engine is allowed to encrypt session data using the decrypted session key to be transmitted to a second device.

11. The system of claim 10, wherein the wrapping key is used to encrypt a group of session keys for a plurality of sessions.

12. The system of claim 10, wherein the secure process is further to transmit, at a series of times and over the system bus, new wrapping keys to be used by the crypto engine to encrypt future session keys.

13. The system of claim 10, wherein the software agent is one of a plurality of protocol-specific software agents executing within a central root of trust on the system to generate session keys to be used for encryption by one of a plurality of protocol-specific inline crypto engines executing on the system.

14. The system of claim 10, wherein the secure processor is further to receive an initial wrapping key and encrypt the wrapping key using the initial wrapping key, the secure processor further to transmit the wrapping key, encrypted using the initial wrapping key, over the system bus to be decrypted by the key unwrap engine and transmitted to the crypto engine, wherein the initial

wrapping key is to be received without encryption during an early execution process of a computing device.

15. The system of claim 10, wherein the key unwrap engine is further to analyze metadata associated with the session key to determine information about the session key, the information indicating at least one of a location to route the session key, an identifier for the session key, or a type of the session key.

16. A key tunnel including one or more processors for receiving session keys transmitted over a system bus of a computing device, the session keys encrypted within a root of trust on the computing device using a wrapping key generated within the root of trust, the one or more processors of the key tunnel to unwrap the session keys within the computing device using the wrapping key and transmit the session keys to one or more protocol-specific crypto engines within the computing device for encrypting session data to be transmitted to one or more remote devices.

17. The key tunnel of claim 16, wherein the one or more processors are further to analyze metadata associated with the session key to determine where to route the session key.

18. The key tunnel of claim 16, wherein the one or more processors are further to analyze a status bit for a received key to determine whether a type of the received key is a session key type or a wrapping key type.

19. The key tunnel of claim 16, wherein the root of trust is to use an initial shared secret to encrypt the wrapping key to be received by the key tunnel, wherein the initial shared secret is one of a plaintext key, an embedded key, a provisioned key shared by two or more fuse blocks, a secret received over a physical private bus, or a secret generated using a specified cryptographic scheme.

20. The key tunnel of claim 16, wherein the key tunnel is comprised in at least one of: a system for performing simulation operations; a system for performing simulation operations to test or validate autonomous machine applications; a system for rendering graphical output; a system for performing deep learning operations; a system implemented using an edge device; a system for generating or presenting virtual reality (VR) content; a system for generating or presenting augmented reality (AR) content; a system for generating or presenting mixed reality (MR) content; a system incorporating one or more Virtual Machines (VMs); a system implemented at least partially in a data center; a system for performing hardware testing using simulation; a system for synthetic data generation; a collaborative content creation platform for 3D assets; or a system implemented at least partially using cloud computing resources.
