

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12407530
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Nickels; Kevin et al.

Automation of user identity using network protocol providing secure granting or revocation of secured access rights

Abstract

The present disclosure involves systems, software, and computer implemented methods for user-controlled access control for user information. One example method includes sending an authentication request to authenticate as a requesting entity to a first decentralized resource directory of a providing entity. An authentication challenge is received, via the connection, from the providing entity, and in response to the authentication request, to store an authentication challenge value for an authentication challenge key in a second decentralized resource directory of the requesting entity. The authentication challenge value for the authentication challenge key is stored in the second decentralized resource directory. An authentication challenge response is sent to the providing entity requesting the providing entity to verify the authentication challenge. An indication is received from the providing entity indicating that the requesting entity is authenticated to the first decentralized resource directory as the requesting entity.

Inventors: Nickels; Kevin (San Jose, CA), Constable; Colin (Mariposa, CA), Vanguri; Jagannadh (Telangana, IN)

Applicant: Atsign, Inc. (San Jose, CA)

Family ID: 79552100

Assignee: Atsign, Inc. (San Jose, CA)

Appl. No.: 18/485932

Filed: October 12, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240039739 A1	Feb. 01, 2024

Related U.S. Application Data

continuation parent-doc US 17781464 US 11849053 WO PCT/US2021/040628 20210707 child-
doc US 18485932
us-provisional-application US 63105755 20201026
us-provisional-application US 63062092 20200806
us-provisional-application US 63049460 20200708

Publication Classification

Int. Cl.: H04L9/32 (20060101); H04L9/08 (20060101)

U.S. Cl.:

CPC H04L9/3271 (20130101); H04L9/0825 (20130101); H04L9/0894 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6343361	12/2001	Nendell et al.	N/A	N/A
8594632	12/2012	Azizi et al.	N/A	N/A
9083531	12/2014	Chenna	N/A	N/A
9258118	12/2015	Roth et al.	N/A	N/A
9571480	12/2016	Asveren	N/A	N/A
11271933	12/2021	Dee	N/A	N/A
11546337	12/2022	Le et al.	N/A	N/A
11849053	12/2022	Nickels et al.	N/A	N/A
2004/0044727	12/2003	Abdelaziz et al.	N/A	N/A
2007/0118879	12/2006	Yeun	N/A	N/A
2008/0091615	12/2007	Choi et al.	N/A	N/A
2015/0178515	12/2014	Cooley et al.	N/A	N/A
2016/0182500	12/2015	Ligatti et al.	N/A	N/A
2017/0257358	12/2016	Ebrahimi et al.	N/A	N/A
2018/0367526	12/2017	Huang et al.	N/A	N/A
2019/0149539	12/2018	Scruby	N/A	N/A
2020/0085408	12/2019	Lemaitre	N/A	H04L 12/1818
2020/0145451	12/2019	Hansen	N/A	N/A
2020/0252392	12/2019	Mullins et al.	N/A	N/A
2021/0058374	12/2020	Smith et al.	N/A	N/A
2021/0110015	12/2020	McCarty et al.	N/A	N/A
2021/0272120	12/2020	Murdoch et al.	N/A	N/A
2021/0288973	12/2020	Dimble et al.	N/A	N/A
2023/0015819	12/2022	Nickels et al.	N/A	N/A
2024/0146533	12/2023	Constable et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
WO 2019/226115	12/2018	WO	N/A
WO 2022/266364	12/2021	WO	N/A

OTHER PUBLICATIONS

International Preliminary Report on Patentability in International Appln. No. PCT/US2021/040628, dated Jan. 19, 2023, 12 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/US2022/033854, mailed on Dec. 14, 2023, 7 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/US2021/040628, mailed on Dec. 2, 2021, 15 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/US2022/033854, mailed on Oct. 25, 2022, 8 pages. cited by applicant

Li et al. "Research on CoAP Resource Directory Based on Blockchain." IEEE 4th Information Technology, Networking, Electronic and Automation Control Conference (ITNEC), May 2020, 1:1730-1735. cited by applicant

Ylonen et al., "Security of interactive and automated access management using Secure Shell (SSH)," National Institute of Standards and Technology, Oct. 2015, 50 pages. cited by applicant

Extended European Search Report in European Appln. No. 22825850.5, mailed on Jun. 27, 2024, 9 pages. cited by applicant

Primary Examiner: Noaman; Bassam A

Attorney, Agent or Firm: Fish & Richardson P.C.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. application Ser. No. 17/781,464, filed Jun. 1, 2022, which is a National Stage Application under 35 USC § 120 to PCT Application Serial No. PCT/US2021/040628, filed Jul. 7, 2021; and published as WO 2022/010978 on Jan. 13, 2022; which claims priority under 35 USC § 119(e) to U.S. Provisional Patent Application Ser. No. 63/049,460, filed on Jul. 8, 2020; and also claims priority under 35 USC § 119(e) to U.S. Provisional Patent Application Ser. No. 63/062,092, filed on Aug. 6, 2020; and also claims priority under 35 USC § 119(e) to U.S. Provisional Patent Application Ser. No. 63/105,755, filed on Oct. 26, 2020, the entire contents each and together are hereby incorporated by reference.

TECHNICAL FIELD

(1) The present disclosure relates to computer-implemented methods, software, and systems for user-controlled access control for user information.

BACKGROUND

(2) Peer-to-peer (P2P) computing or networking is a distributed application architecture that can include partitioning of tasks between peer computing devices. Peers can form a peer-to-peer network of nodes. In a P2P environment, peers can make a portion of their resources, data, and/or metadata, such as processing power, disk storage or network bandwidth, directly available to other network participants, without requiring central coordination. Peers can be both suppliers and consumers of resources, in contrast to a traditional client-server model in which servers supply resources and clients consume resources.

SUMMARY

(3) The present disclosure involves systems, software, and computer implemented methods for user-controlled access control for user information.

(4) An example method includes: receiving, via a first connection and from a first requesting entity, an authentication request to authenticate the first requesting entity to a first decentralized resource directory of a providing entity, wherein the authentication request identifies the first requesting entity; generating, in response to the authentication request, an authentication challenge value and an authentication challenge key; providing, via the first connection and in response to the authentication request, an authentication challenge to the first requesting entity for the first requesting entity to store the authentication challenge value for the authentication challenge key in a second decentralized resource directory of the first requesting entity; receiving, via the first connection, a confirmation from the first requesting entity that the authentication challenge value has been stored for the authentication challenge key in the second decentralized resource directory of the first requesting entity; establishing a second connection with the second decentralized resource directory of the first requesting entity; sending, via the second connection, a first lookup request for a value stored for the authentication challenge key in the second decentralized resource directory; receiving, via the second connection and from the second decentralized resource directory, a first response value in response to the first lookup request; comparing the first response value to the authentication challenge value to determine whether the first response value matches the authentication challenge value; and in response to determining that the first response value matches the authentication challenge value, responding to the authentication request, via the first connection, with an indication that the first requesting entity is authenticated to the first decentralized resource directory.

(5) Implementations may include one or more of the following features. In response to determining that the first response value does not match the authentication challenge value, the first connection can be dropped. A second lookup request for a value for a key stored in the first decentralized resource directory can be received, via the first connection and from the first requesting entity, while the first requesting entity is authenticated to the first decentralized resource directory. A determination can be made that the first requesting entity has been provided access to a first stored value for the key. The first stored value for the key can be retrieved and the first stored value can be provided as a second response value to the first requesting entity, in response to the second lookup request. A third lookup request can be received, via an unauthenticated third connection and from an unauthenticated second requesting entity, for a value for the key stored in the first decentralized resource directory. A publicly accessible value for the key can be retrieved and the publicly accessible value for the key can be provided as a third response value, to the second requesting entity, in response to the third lookup request. The first stored value for the key provided to the authenticated first requesting entity can be different from the publicly accessible value for the key. A fourth lookup request for a value for the key stored in the first decentralized resource directory can be received, via a third connection and from a third requesting entity while the third requesting entity is authenticated to the first decentralized resource directory. A determination can be made that the third requesting entity has been provided access to a third stored value for the key. The third stored value can be different from the first stored value and the publicly accessible value for the key. The third stored value for the key can be retrieved and provided as a third response value to the third requesting entity, in response to the fourth lookup request. A connection can be made, before sending the second lookup request, to a namespace directory to request connection information for the second decentralized resource directory. The connection information for the second decentralized resource directory can be used to establish the second connection. The authentication request and the second lookup request can be received from a user device associated with the first requesting entity. The authentication request and the second lookup request can be received from a server that services requests for the second decentralized resource directory. The

authentication challenge value can be a randomly-generated identifier.

(6) Another example method includes: sending, via a first connection, an authentication request to authenticate as a requesting entity to a first decentralized resource directory of a providing entity, wherein the authentication request identifies the requesting entity; receiving, via the first connection, from the providing entity, and in response to the authentication request, an authentication challenge to store an authentication challenge value for a authentication challenge key in a second decentralized resource directory of the requesting entity; storing, in the second decentralized resource directory, the authentication challenge value for the authentication challenge key; sending an authentication challenge response to the providing entity requesting the providing entity to verify the authentication challenge; and receiving an indication from the providing entity that the requesting entity is authenticated to the first decentralized resource directory as the requesting entity.

(7) Implementations may include one or more of the following features. A first lookup request can be sent to the first decentralized resource directory for a value for a key stored in the first decentralized resource directory, while authenticated to the first decentralized resource directory and via the first connection. A first value for the key can be received, in response to the first lookup request, a first value for the key. A second lookup request can be sent, as an unauthenticated requesting entity, via a second connection, to the first decentralized resource directory for the value for the key stored in the first decentralized resource directory. A second value for the key can be received, in response to the second lookup request, with the second value being different from the first value. A connection can be made, before sending the authentication request, to a namespace directory to request connection information for the first decentralized resource directory. The connection information can be used by the second decentralized resource directory to establish the first connection. The authentication challenge value can be a randomly-generated identifier generated by the providing entity.

(8) An example system includes: a user device of a first user; a first protocol server associated with the first user comprising a first decentralized resource directory; a second protocol server associated with a second user comprising a second decentralized resource directory; and a namespace directory. The user device is configured to send a first lookup request to the namespace directory for first connection information for the first protocol server. The namespace directory is configured to receive the first lookup request from the user device, determine the first connection information for the first protocol server, and provide the first connection information to the user device. The user device is further configured to use the first connection information to send a first authentication request to the first protocol server. The first protocol server is configured to receive the first authentication request, generate a key challenge, and send the key challenge to the user device in response to the first authentication request. the user device is further configured to generate a first key challenge response and send the first key challenge response to the first protocol server. The first protocol server is further configured to verify the first key challenge response and send a first authentication success indication to the user device. The user device is further configured to send a second lookup request to the first protocol server for a value of a requested key stored in the second decentralized resource directory. The first protocol server is further configured to receive the second lookup request and send a third lookup request to the namespace directory for second connection information for the second protocol server. The namespace directory is further configured to receive the third lookup request, determine the second connection information for the second protocol server, and send the second connection information to the first protocol server. The first protocol server is further configured to receive the second connection information and use the second connection information to send a second authentication request to the second protocol server. The second protocol server is configured to receive the second authentication request, generate an authentication challenge, and send the authentication challenge to the first protocol server. The first protocol server is further configured to receive the

authentication challenge, store an authentication challenge value in the first decentralized resource directory for an authentication challenge key, and send a confirmation for the authentication challenge to the second protocol server. The second protocol server is further configured to receive the confirmation for the authentication challenge and send a fourth lookup request for third connection information, for the first protocol server, to the namespace directory. The namespace directory is further configured to receive the fourth lookup request, determine the third connection information for the first protocol server, and send the third connection information to the second protocol server. The second protocol server is further configured to use the third connection information to send a fifth lookup request to the first protocol server for a stored value for the authentication challenge key, receive the stored value for the authentication challenge key, verify the stored value for the authentication challenge key, and send a second authentication success indication to the first protocol server in response to a successful verification of the authentication challenge key that indicates successful authentication to the second protocol server as the first user. The first protocol server is further configured to receive the second authentication success indication and forward the second lookup request, to the second protocol server, on behalf of the user device. The second protocol server is further configured to receive the second lookup request, determine that the first user has access to the requested key, retrieve the value for the requested key, and send the value for the requested key to the first protocol server. The first protocol server is further configured to receive the value for the requested key and forward the value for the requested key to the user device. the user device is further configured to receive the value for the requested key.

(9) Implementations may include one or more of the following features. The system can include a registration server. The user device can be configured to receive a shared secret from the registration server during a registration process and store the shared secret. The first protocol server can be configured to receive the shared secret from the registration server and store the shared secret. The shared secret can be used during a bootstrap process performed by the user device and the first protocol server. The shared secret can be used in a cryptographic response authentication mechanism. The user device can be further configured to generate, after the bootstrap process, a private key and a public key corresponding to the private key, store the private key on the user device, and send the public key to the first protocol server. The first protocol server can be further configured to receive the public key and store the public key at the first protocol server. The first protocol server can be further configured to use the public key to generate the key challenge. The user device can be further configured to use the private key to generate the first key challenge response. The first protocol server can be further configured to use the public key when verifying the first key challenge response.

(10) Another example method includes: receiving, via a first connection and from a first requesting entity, an authentication request to authenticate the first requesting entity to a first decentralized resource directory of a providing entity, wherein the authentication request identifies the first requesting entity; generating, in response to the authentication request, an authentication challenge value and an authentication challenge key; providing, via the first connection and in response to the authentication request, an authentication challenge to the first requesting entity for the first requesting entity to store the authentication challenge value for the authentication challenge key in a second decentralized resource directory of the first requesting entity; receiving, via the first connection, a confirmation from the first requesting entity that the authentication challenge value has been stored for the authentication challenge key in the second decentralized resource directory of the first requesting entity; establishing a second connection with the second decentralized resource directory of the first requesting entity; sending, via the second connection, a first lookup request for a value stored for the authentication challenge key in the second decentralized resource directory; receiving, via the second connection and from the second decentralized resource directory, a first response value in response to the first lookup request; comparing the first response

value to the authentication challenge value to determine whether the first response value matches the authentication challenge value; and in response to determining that the first response value matches the authentication challenge value, responding to the authentication request, via the first connection, with an indication that the first requesting entity is authenticated to the first decentralized resource directory.

(11) While generally described as computer-implemented software embodied on tangible media that processes and transforms the respective data, some or all of the aspects may be computer-implemented methods or further included in respective systems or other devices for performing this described functionality. The details of these and other aspects and embodiments of the present disclosure are set forth in the accompanying drawings and the description below. Other features, objects, and advantages of the disclosure will be apparent from the description and drawings, and from the claims.

Description

DESCRIPTION OF DRAWINGS

- (1) FIG. 1 is a block diagram illustrating an example system for user-controlled access control for user information.
- (2) FIG. 2A is a block diagram illustrating an example system for user-controlled access control for user information.
- (3) FIG. 2B is a block diagram illustrating an example system for a public lookup transaction.
- (4) FIG. 3 is a block diagram illustrating an example system for a private lookup transaction.
- (5) FIGS. 4A and 4B are a combined swim lane diagram illustrating an example method for a private lookup transaction.
- (6) FIG. 5 illustrates example keys of values stored in a user's personal directory.
- (7) FIG. 6A-C illustrates example uses of a scan verb.
- (8) FIG. 7 illustrates a session that includes use of an example profile and an example addressing scheme.
- (9) FIG. 8 is a block diagram illustrating an example system for a private lookup transaction.
- (10) FIGS. 9A-9B are a combined swim lane diagram illustrating an example method for a private lookup transaction.
- (11) FIG. 10 is a block diagram illustrating an example system for configuring a device for protocol exchanges.
- (12) FIGS. 11A-11B are a combined swim lane diagram illustrating an example method for configuring a device for protocol exchanges.
- (13) FIG. 12 is a flowchart of an example method for receiving and responding to a protocol request for user information.
- (14) FIG. 13 is a flowchart of an example method for providing a protocol request for user information.
- (15) Onboarding and Initialization Options
- (16) FIG. 14 is a flow chart of an example onboarding process.
- (17) FIG. 15 is a flow chart of an example onboarding process for onboarding a second application.
- (18) FIG. 16 is a flow chart of an example onboarding process using a separate private key per application.
- (19) FIG. 17 is a flow chart of an example onboarding process for a second application using a separate private key per application.
- (20) FIG. 18 is a flow chart of an example onboarding process using a server-based private key.
- (21) FIG. 19 is a flow chart of an example onboarding process for a second application using a server-based private key.

- (22) FIG. 20A illustrates an example system for onboarding a device/application.
- (23) FIG. 20B illustrates an example system for onboarding a device/application using a registrar.
- (24) FIG. 20C is a flow chart of an example onboarding process for onboarding a device/application using a registrar.
- (25) FIG. 21 is a flowchart of an example method for lifecycle management of a secondary server.
- (26) FIG. 22 is a flowchart of an example method for data expiration management.
- (27) FIG. 23 illustrates an example system for saving a profile using encryption.
- (28) FIG. 24 illustrates a system for @sign activation.
- (29) FIG. 25 illustrates @sign namespaces.
- (30) Streaming and File Transfers
- (31) FIG. 26 illustrates a system that illustrates stages of protocol-enabled file sharing.
- (32) FIG. 27 is a flowchart of example streaming method.
- (33) FIG. 28 is a flowchart of example streaming method using a server hop.
- (34) FIG. 29 is a flowchart of example streaming method using a token.
- (35) FIG. 30 is a flowchart of example streaming method.
- (36) FIG. 31 is a flowchart of example streaming method using a proof of life approach.
- (37) FIG. 32 is a flowchart of example streaming method using an intermediate streaming server.
- (38) Data Access, Sharing and Encryption
- (39) FIG. 33 illustrates an example system for use of encryption keys.
- (40) FIG. 34 illustrates an example system for use of encryption keys.
- (41) FIG. 35 illustrates an example system for key value retrieval.
- (42) FIG. 36 illustrates approaches for spam protection.
- (43) Local and Secondary Server Synchronization
- (44) FIG. 37 illustrates example synchronization operations for synchronizing local and cloud persistence.
- (45) FIGS. 38A-C illustrate a synchronization scenario in which a local storage is ahead (update wise) of storage on a secondary server.
- (46) FIG. 39 illustrates a synchronization scenario in which a local storage is behind (update wise) storage on a secondary server.
- (47) FIG. 40 illustrates a synchronization scenario in which a local storage becomes synchronized with a secondary server after data is deleted from the secondary server.
- (48) FIG. 41 illustrates a synchronization scenario in which a local storage becomes synchronized with a secondary server after data is deleted from the secondary server and added to the local storage.
- (49) FIGS. 42A-B illustrate a synchronization scenario in which a local storage becomes synchronized with a secondary server after data is deleted from the local storage and added to the secondary server.
- (50) FIG. 43 illustrates an example system for synchronizing keys across user devices and a secondary server.
- (51) FIGS. 44A-B illustrate an example system for synchronizing keys across applications and a secondary server.
- (52) FIG. 45 illustrates an example sync response message.
- (53) Design
- (54) FIG. 46 illustrates example entities for verb handling.
- (55) FIG. 47 illustrates an example class diagram.
- (56) FIG. 48 illustrates client design criteria goals
- (57) Notifications
- (58) FIG. 49 illustrates an example system for an accept notification.
- (59) FIG. 50 illustrates an example system for an update notification.
- (60) FIG. 51 illustrates a peer to peer network for group notifications.

- (61) FIG. 52 illustrates an example reflector network.
- (62) FIG. 53 illustrates information describing a notification verb.
- (63) FIG. 54 illustrates an example system for notifications.
- (64) FIG. 55 illustrates an example system for notifications.
- (65) Example Applications
- (66) FIG. 56 illustrates an example system for P2P applications.
- (67) FIGS. 57A-B are conceptual diagrams that illustrate an example P2P application for sharing data between an executive and an executive assistant.
- (68) FIG. 58 illustrates an example configuration user interface for a P2P data sharing application.
- (69) FIG. 59 illustrates an example executive assistant list user interface for a P2P data sharing application.
- (70) FIG. 60 illustrates an example executive assistant permissions user interface for a P2P data sharing application.
- (71) FIG. 61 illustrates an example executive list user interface for a P2P data sharing application.
- (72) FIG. 62 illustrates an example data access user interface for a P2P data sharing application.
- (73) FIG. 63 is a block diagram illustrating an example system for execution of a P2P data sharing application.
- (74) FIG. 64 illustrates a user interface for an improved contacts application.
- (75) FIG. 65 illustrates an example invitation user interface for the improved contacts application.
- (76) FIG. 66 illustrates an example relationship user interface for displaying relationship information for the improved contacts application.
- (77) FIG. 67 illustrates an example persona selection user interface.
- (78) FIG. 68 illustrates an example user interface for viewing personas.
- (79) FIG. 69 illustrates an example persona user interface.
- (80) FIG. 70 illustrates an example relationships user interface.
- (81) FIG. 71 illustrates an example group creation user interface.
- (82) FIG. 72 illustrates an example group creation user interface.
- (83) FIG. 73 illustrates an example group creation user interface
- (84) FIG. 74 illustrates an example system for attestations.

DETAILED DESCRIPTION

(85) As the Internet has grown into every facet of modern life, some of the most problematic issues people face are managing their identities and securing their data. Additionally, website and application developers may have to manage the identity of people thereby assuming an associated risk of managing data of other parties and incurring associated hosting and infrastructure costs. Applications that are monetized through advertising that includes tracking and identifying users can further result in various privacy concerns.

(86) To solve these and other problems, an improved and novel Internet Protocol (IP)-based communication protocol and system can be used to provide robust privacy controls for secure exchanges of information between entities. In particular, the solution provides protocol users with the ability to personally maintain and manage sharing of their personal information and data without reliance on a third-party service or application's settings and restrictions. The system includes a namespace directory, decentralized resource directories with secure key/value persistence, and a protocol (which can be referred to as the “@protocol”) for communication between entities. Information, including both data and services, can be securely and privately exchanged over the protocol.

(87) As used herein, “users” or “protocol users” refer to users of the @protocol. Each protocol user can be a person, and it follows that multiple users can be “people.” A respective person can use the protocol, with any suitable device, including their own personal device, which can be referred to a “user device,” or “client device.” In general, a computing device that participates in the protocol can be associated with a participating entity, which can be a person, a corporate or other broader

entity, or a finer-grained entity such as a specific item or “thing” (such as an IoT (Internet of Things)) device. Each entity that uses the protocol can be represented by a handle, which can be referred to as an “@sign”. Example @signs can include @alice and @bob. An @sign can be combined with a service. For example, sales@amco and myfurnace@alice refer to a sales service for the amco @sign and a myfurnace service for the @alice @sign, respectively. For a corporate or other entity, various handles can be supported, such as a handle for the entity itself (e.g., @amco), a handle for a general role for the entity (e.g., sales@amco), and a handle for specific users associated with the entity (e.g., alice.smith@amco).

(88) The namespace directory can be used for storing a list of unique identifiers (e.g., “@signs”) of protocol users together with routing information for the identifiers. Routing information can include, for example, a DNS (Domain Name System) location and IP port number of individual decentralized resource directory servers (which can be referred to as @sign servers) to which protocol requests can be routed.

(89) A response to a data or service request can be provided by a decentralized resource directory server based on a secure key/value persistence of the decentralized resource directory. A request can include an indication of a requested data item or service, as well as a target entity of the request. For example, a service-related request can have a syntax of a form of “service@entity.” Processing of a protocol request can include operations to ensure that both entities in an exchange are known to each other and that the request is permission checked before an appropriate response is determined and returned. Response information returned for a service@entity request can include or refer to (e.g., by value or by reference) information stored in a secure key/value persistence for the entity.

(90) The @protocol can allow for a person or an entity to configure the person or entity's own data records to be either publicly accessible (e.g., available to anyone who is aware of a corresponding key) or to be private (e.g., with limited, controlled access) and require validation of a requesting entity for access, where only people or entities specified by the person or entity who owns the data records can access the data, or corresponding portions if limited by the person or entity. As described in more detail below, some keys for public data can be discovered/exposed using a scan while other keys may not be discoverable by a scan. Access control for private data can be hard coded into the system rather than as an overlay data structure. To access private data, a requesting entity can be required to first prove, using the @protocol, a claimed identity, so that access rights can be determined. Accordingly, a response to a private data request can be based on a validated identifier of the requestor, and responses can vary between different requesting entities based on their respective level of approved access. For example, a person can configure permissions so that a query sent to the person for the person's location can result in different responses of “USA,” “Minnesota,” or an actual physical address, if sent from an unknown sender, a casual acquaintance, or a close friend, respectively, based on respective grants previously granted to respective entities. The person can also configure permissions for the location data so that no response is given to requests received from unknown senders.

(91) The features of the protocol can form a unique Internet permissioning system for securely controlling access to data and services. Permissioned access can be controlled strictly by an owner of a specific private namespace granting access to data or service(s) of the owner. Access grants can be subject to revocation by the owner.

(92) The permissioning system can provide various advantages. Several advantages relate to an enhanced level of control for people of their own personal information. For example, people are able to specify their own terms and conditions for the use of their data. Specified terms can be defined in a data-use license and/or terms of service document, for example. A person can specify in a data-use license that a particular data item can be looked up but not stored or replicated, for instance.

(93) The @protocol is designed to only allow connections to occur when permission has been

granted by the owner of each respective @sign handle. The response to a query, including whether to respond at all, as well as the payload of a response, is completely under the control of the owner of the @sign, which can allow for fine- or coarse-grained control of responses, which can include different responses to a same query to different users or entities. As a particular example, the protocol can enable a person, such as a person or user associated with the @sign of “@alice,” to share information, such as email, websites, credit card numbers, location data, or preferences, to other people or entities, such as a person “@bob,” under the complete control of and according to the personally-defined rules and permission of @alice alone.

(94) Providing each person or entity control of their own information can improve user experiences in a variety of ways. For example, the described method of directly asking Internet users for their information can allow for new application experiences providing less annoyance for both the requester and the user to whom the request is sent. The @protocol, with a set of predefined verbs, which can be different in (e.g., extended for) particular implementations, and can provide an ease of use and intuitiveness for protocol users. The system can provide people, entities, and even physical items (e.g., “things”)—or authorized users or persons associated therewith—a single location to manage digital information and a simple mechanism to allow others to request that information. Each protocol user, entity, or physical item can have and share a respective @sign, which can be used by others for future data requests. A simple @sign can be shared (e.g., on or as a business card, via text, email, or verbally, etc.), rather than separate items of information that may become outdated at some point. Sharing an @sign (e.g., @alice or sales@acme) can enable the contacting of Alice @alice or the sales role at AcmeCo, respectively, even if Alice's actual contact details change or if the actual person or people in the sales role at Acme company change.

(95) Generally, the protocol provides a hierarchical and decentralized permissioning system for establishing security and trust as a core part of the Internet, whether between people or between people and entities (e.g., brands). Respective @sign servers can participate in the @protocol without being surveilled or managed by a central entity. Additionally, end-to-end encrypted P2P applications that exchange data using the @protocol can be used without being subject to centralized surveillance, since they do not communicate through centralized and shared backend servers. Rather, protocol-enabled P2P applications can communicate directly with other applications at the edge of the decentralized network after a connection between peers is established.

(96) As another example, data managed by the system can be housed, for instance, in billions of independent @sign servers (which can be implemented as server devices and/or user devices), with each server storing data in a secure and encrypted manner. In contrast to a large database housed on a centralized system that may be vulnerable to a single security attack, any security attack that might be undertaken against the decentralized permissioning system would have to be performed separately against each @sign server, which would be cost and time prohibitive. Decentralized systems can be more trustworthy than a centralized system that has an inherently higher risk (and perhaps history) of data breaches for centrally-stored data. Historical security issues with centralized systems generally result in erosion of trust in the centralized system.

(97) The @protocol is designed as a low-level, real-time internet permissions protocol. The permissions-based protocol is low-level in that it can prevent a transfer of secured data until permission is verified. Permissions are checked prior to any data transfer, and data is only transferred in response to a successful verification, thereby providing a near real time response that ensures current data, in some cases, without—or with reduced—caching. Application developers (and in some cases, application users) can specify caching parameters for an application, such as whether caching of data is intended or preferred, for instance.

(98) The system supporting the @protocol is based on a global-scale architecture that is dependable as the system scales. Distributed @sign servers and an addressing scheme can support at least one hundred billion people, entities, and things, for example. Dependability is achieved based on

architectural designs for confidentiality, integrity, and availability.

(99) The open-source @protocol has a well-defined architecture and a reference implementation codebase that can be published to third party reviewers to facilitate code review and extensive security testing. The architecture is grounded in fundamental tenets of 1) distributing data when possible and centralizing data when necessary, 2) requiring permission for data access, and 3) empowering ownership of data. The @protocol provides a well-defined set of functionality that is extensible, enabling developers to add useful extensions and take advantage of the security and protection offered by the solution.

(100) Application developers can beneficially leverage various aspects of the @protocol. With protocol-enabled applications, users own and control their own data, so applications can be automatically compliant with GDPR (General Data Protection Regulation), CCPA (California Consumer Privacy Act), and other current and future privacy regulations. With personal data used by an application being stored with the person, the person, and not an application provider, stores keys and other sensitive information, such as personally identifiable information (PII). There is no longer a reason for the developer to add warnings to forms or other areas of an application or site where the user is entering data, as may be required for other applications and prior solutions. Automatic privacy compliance can result in a more empowering, reassuring experience for the user and less effort for a developer, as well as an inherent security to the system from end to end.

(101) User management, particularly of personal information, therefore, can be minimal (or non-existent) for an application developer, and can result in applications not needing an expensive backend infrastructure. Building an infrastructure to house and back up customer data can be a costly and time consuming part of application development. An application being developed to use the @protocol can save costs and time in getting the application to market, as compared to approaches that need development of a backend infrastructure, including security and storage. Fully decentralized P2P apps that have data stored on users' individual @servers may not require any backend infrastructure, for example, which can save substantial costs for the developer.

(102) Liability for data breaches may no longer exist for an application when the application developer does not hold any of the user's data. If the developer chooses to have personal data hosted on personal servers, for example, the individual people can become responsible for their own data, with the developer using best practices and the described @protocol to request and receive access to the personal data. Developers therefore, may avoid a liability for sensitive personal data, and can limit the amount of collected and managed data held by the application.

(103) Additionally, data management can be simplified, both for the developer and for the user. For example, password retrieval and out-of-date email addresses, if managed by the application developer, can create a large support burden for application providers. By using the @protocol, developers no longer have to develop functionality for and support password, login, and other personal data management. Once a user has an @ handle, they do not need to create a username or password for @protocol-compatible sites or applications—instead, they can simply use their @handle, and can verify themselves via the @protocol solution as described herein. Accordingly, the user can get started more easily with a new protocol-enabled application, resulting in a more satisfying user experience, which can benefit the developer as well as the user.

(104) FIG. 1 is a block diagram illustrating an example system **100** for user-controlled access control for user information. Specifically, the illustrated system **100** includes or is communicably coupled with various components that are connected by one or more public network(s) **101**. For example, the system **100** includes a governing body **102**, a protocol provider server **103**, application developers **104**, root servers **105**, backup servers **106**, user devices **107**, an application store **108**, protocol servers **109** (e.g., @sign servers), and registrar servers **110**. Additionally, in some implementations, the system **100** includes enterprise user devices **111** that might not be directly connected to the public network(s) **101**, but are, for example, on a local area network **112**, behind a firewall in an enterprise, and connected to an enterprise microservice server **113** that is

connected to the local area network **112** and the public network **101** that can do protocol exchanges on behalf of the enterprise user devices **111**. Although shown separately, in some implementations, functionality of two or more systems or servers may be provided by a single system or server. In some implementations, the functionality of one illustrated system or server may be provided by multiple systems or servers.

(105) The system **100** provides an ecosystem of compliant applications and services for granting or denying access rights in real time. The protocol, as implemented in the system **100**, can create a network of trust with structured data on a global scale. The protocol can be built upon protocols of TCP/IP (Transmission Control Protocol/Internet Protocol), DNS, HTTP (HyperText Transfer Protocol) (e.g., with future browser enhancements) and SSL/TLS (Secure Sockets Layer/Transport Layer Security). Use of the protocol can result in secure sharing of data and services between @handle owners and applications as specified in @sign servers, if explicitly allowed by both parties of a transaction. The protocol can be used for a large number of use cases that can benefit from user controlled privacy access. The protocol supports both a public lookup (in which the party querying data does not have to prove who they are) and a private lookup that confirms the @sign of the entity asking for information and allows the owner of the data to decide what information to share. Public lookups in general can be used to confirm the existence of an @sign or access specific information that may be set to public access.

(106) The protocol can be provided/developed by a protocol provider (shown interfacing with the system **100** by the protocol provider server **103**). Although a single protocol provider server **103** is shown, the protocol provider can be associated with various types and numbers of servers. And while shown as a single entity, the protocol provider can include or encompass different entities, such as a non-profit organization/foundation tasked with promoting and supporting the protocol and/or a company that provides various services and policies for the protocol.

(107) For example, the protocol provider server **103** can provide a source code distribution **114** and a documentation distribution **116**, include an application certifier **118** that generates application certification information **119**, and include a server configuration component **120** for configuring the root servers **105**. As another example, the protocol provider server **103** can provide secondary user server infrastructure and services, using user containers **122** and user microservices **124**. The components of the protocol provider server **103** are described in more detail below.

(108) In some implementations, the protocol provider server **103** provides governance of the protocol namespace. In other implementations and/or at different points in time, governance can be provided by the governing body **102**. The governing body **102** can be ICANN (Internet Corporation for Assigned Names and Numbers) or some other entity. The governing body **102** (and/or the protocol provider) can include a namespace governor **126** that governs the @sign namespace, including arbitration of namespace conflicts. The governing body **102** (and/or the protocol provider) can also include a registrar accreditor **128** for accreditation of registrars that can sell and maintain @signs and secondary server configurations.

(109) In some implementations, the protocol provider (e.g., as the protocol provider server **103** or other servers or components) can act as a registrar. Other third party registrars can be supported in some implementations. For example, the registrar servers **110** can provide a registration service **130** that maintains @sign data **132** by performing the registering, provisioning, transferring, renewal, and deletion of @sign handles for users.

(110) An @sign, such as @alice, is a unique identifier assigned to a user of the protocol. An @sign can be referred to as a handle, or an @sign handle, and can be a string of characters (e.g., Unicode characters) uniquely identifying a user and that is registered with a registrar. An @sign can serve as a personal or corporate brand. An expanded Unicode character set can support a wide variety of characters for creative @sign handle creation. Within the system **100**, an @sign serves as an address that is resolvable by the root servers **105** to a secondary server endpoint, as described in more detail below.

(111) The root servers **105** are an implementation of a global namespace directory for the system **100**. The root servers **105** can store a list of unique identifiers (@signs) along with information for how to reach an @sign server that handles requests for a given @sign. For example, the root servers **105** can store and maintain mappings **137** that map @signs to location information for Internet-accessible user protocol microservices. Location information can be a DNS address plus a port number. The user protocol microservices, referred to herein as @sign servers, @servers, or protocol servers, implement individual, distributed, decentralized resource directories to which protocol requests are routed.

(112) The root servers **105** can provide redirection to the @server for a particular @sign. For example, the root servers **105** can offer a centralized lookup service **140** that provides access to a single @sign namespace. Unlike DNS, the @sign namespace can be implemented without Top Level Domains (TLDs) or subdomains. The lookup service **140** can be intentionally minimal for efficiency, offering just two items of information: 1) whether a particular @sign exists; and 2) if the @sign exists, how the server for the @sign can be reached.

(113) The lookup service **140** can support an @exit verb which can disconnect a connection with a requester. All other inputs to the lookup service **140** can be considered to be @sign lookup requests. Location information for @sign servers can be considered public information so the lookup service **140** can be configured to accept requests without requiring authentication. Users can first query the lookup service **140** for an Internet-addressable location of an @sign server before sending protocol requests to that @sign server.

(114) In response to a lookup query, the lookup service **140** can respond with a null value (or simply ignore a request) if a requested @sign does not exist. If the requested @sign does exist, then the lookup service **140** can return the DNS name or address (e.g., a Fully Qualified Domain Name (FQDN)) of the @server and the IP port number for the requested @sign. If a DOS (Denial Of Service) attack is detected, the lookup service **140** can reduce response time for replying to requests and/or terminate connections or restart the lookup service **140**. Other approaches for dealing with DOS attacks can be used. For example, the lookup service **140** can provide the IP address or other identifying information of an entity suspected of a DOS or other attack to upstream (or other) networks or entities, to either slow or delete DOS or other malicious traffic.

(115) TCP ports can be used for scaling @sign server addresses. A TCP port can be used, in general, to identify a service. Port 80 is generally used to identify a web service, for example. Port numbers can generally be a value from 0 to 65535, with only a portion of those values reserved for traditional services. For example, IP version four (IPV4) and IP version six (IPV6) can support port numbers. By using otherwise unassigned or unused port numbers to identify a particular user microservice or @sign server, each unique IP address can be mapped to approximately 65,000 unique microservices, by combining the IP address with a port number. Port numbers thus can allow for scalability at a global scale, since an IP address plus port number scheme can support billions of uniquely-identified microservices.

(116) In general, the root servers **105** can be designed to be cost effectively scaled to handle trillions of requests for billions of records. Root servers **105** can be architected to just provide the single lookup service **140** for the mappings **137**, e.g., to include a minimum amount of information for fulfilling lookup requests. In-memory databases and other performance-improving approaches can be used. For example, multiple read-only copies or replicas (e.g., shards) of the mappings **137** can be used, and a given request can be hashed to identify a particular subset database in which to lookup a given @sign. For instance, a first mappings database can hold @signs starting with the letters “a” through “m” and a second mappings database can hold @signs starting the letters “n” through “z.” As another example, DNS records for multiple root servers **105** can be used to share an overall load across the multiple root servers **105**. For instance, traffic can be routed to a nearest and/or fastest root server **105**, or a root server **105** with a lightest current load.

(117) As yet another example, the root servers **105** can be configured so that read-only replicas

serve the request load and a main root server **105** serves as the master. The main root server **105** can have a preconfigured DNS address (e.g., root.at-sign.org). At least some of the root servers **105**, including the master root server, can be configured, operated, and/or provided by the protocol provider. The protocol provider can manage a cluster of root servers, for example. The master root server and/or the protocol provider can maintain IP addresses of the various root servers. In some implementations, third party provider(s) can provide at least some of the read-only replica root servers. In some implementations, a third party representing a particular organization can deploy their own root server(s) for their organization, and maintain internal addressing for their own namespace that may also bridge to the global @sign namespace.

(118) Referring again to the registrar servers **110**, the @sign data **132** maintained by the registrar servers **110** can include information for user provisioned @signs. The @sign data **132** can also include information for reserved @signs, such as for countries, states, or @signs that correspond to registered trademarks. When a user deletes (e.g., gives up) an @sign, the registration service **130** can mark the @sign as available to other users or entities. Similarly, the registration service **130** can periodically determine @signs that have not been renewed, and identify the unrenewed @signs as available. Renewal dates can also be included in the @sign data **132**. The registration service **130** can periodically determine, in general, which free and/or for-fee @signs are available.

(119) The registrar servers **110** can maintain user contact data **134** for sending renewal reminders or other notices. Financial data **136** can include pricing data for vanity @signs and payment records for @sign purchase transactions. Registrars can store a minimum amount of personally identifiable information (PII), if any. In some implementations, registrars can provide some @sign handles for free and charge a fee for certain (e.g., desirable/vanity) @sign handles.

(120) The registration service **130** can provide a registration application **138** to user devices **107**, for example. The registration application **138** can be provided as an application to install on a given user device **107** or as a web site accessed by the user device **107**. The registration application **138** can be used by users to provision @signs. For example, a user can use the registration application **138** to search for available @signs, view a list of available @signs, view a list of suggested @signs, check if a format of a desired @sign is valid, and view which vanity @signs are available and at what cost. Once an available @sign has been selected (or entered), the user can use the registration application **138** to register the @sign.

(121) Registration can include providing payment to secure a particular @sign, such as for premium or otherwise desired or likely-to-be desired @signs, which may be associated with a cost. Payment can include, for example, an agreement to and payment of a yearly subscription fee. Payment information and details provided through the registration application **138** can be stored in the financial data **136**.

(122) Management and maintenance of @signs can be performed using the registration application **138**. For example, users can use the registration application **138** to receive alerts about upcoming @sign renewals (as generated by the registration service **130**), pay for a renewal, view a dashboard of assigned @signs and renewal dates, or delete (e.g., give up) an @sign. As another example, the user can use the registration application **138** to select an @sign to sell, set up advertising for the for-sale @sign (to be displayed to other potential users on a registration website, for example), or to transfer an @sign to another entity.

(123) The registration application **138** can also be used to bind (or rebind) an @sign to server information for an @sign server. For example, @sign registration can include mapping the registered @sign to server information for an @sign server that can handle protocol exchanges for the @sign. The server information can be a single IP address (and possible port number) or a FQDN and (possible) port number. Using an IP address or a FQDN/port can have advantages and disadvantages.

(124) For example, if a DNS address is used then the IP address mapped to the DNS address can be changed without changing a root entry. However, DNS can cache data meaning migration to a new

IP address for an @sign could take hours or possibly days due to the distributed caching nature of DNS. If an IP address is used as the server information for an @sign an IP address update can be completed more quickly, since protocol guidelines can specify that resolvers or applications may cache IP addresses, for example, for up to 10 minutes. However, if an IP address is mapped directly to an @sign, the IP address for the @sign cannot be changed without contacting the registrar to request an update to a root entry. In some implementations, DNS is preferred over use of IP addresses.

(125) The registrar server **110** can store @sign data for registered and bound @signs in the @sign data **132**. The registrar server **110** can send a request to the root server **105** to create a root entry that maps the @sign to the server information for the @sign. In some implementations, @sign to server mapping information is only held in the root servers **105** (and not in the @sign data **132** maintained by the registrar).

(126) Server information for @sign servers can refer to different types of devices and services. Essentially, an @sign server process can run on any device for which an IP address (and possible port number) given by the root server **105** is an Internet-routable address. For example, the protocol server **109** can be a dedicated (e.g., always-on) protocol-enabled device that can handle requests for a particular @sign. The protocol server **109** can execute a user microservice **142** for a particular @sign handle, providing access (based on assigned access rights) to a personal directory **144** that stores information for the @sign handle. The protocol server **109** can be a server device (managed by a third party), a user device (e.g., personal server) that is continuously connected to the network **101**, an IoT (Internet of Things) device, or some other protocol-enabled, network-connected device.

(127) As another example, a protocol service can run on a user device **107** (e.g., as implemented by a settings application **145** or another process). Since some user devices may be occasionally offline, the user device **107** can be mapped to a secondary server. The secondary server can be the protocol server **109** or can be a cloud-based user microservice provided, for example by the registrar, the protocol provider, or another entity. Third party providers can agree to comply with protocol guidelines before being allowed to act as a cloud provider for secondary server infrastructure.

(128) As an example of cloud infrastructure, the protocol provider server **103** can provide the user microservices **124** as @sign server processes for users for access to user-specific user containers **122**. Similarly, additionally, and/or alternatively, the registrar server **110** can provide user microservices **146** as @sign server processes for users for access to user-specific user containers **148**. User containers **122** or **148** may be associated with TLS certificates. The TLS certificates may be associated with, for example, DNS information for the user microservice **124** or **146**. The TLS certificates can be used to provide authentication of both server and client connections to the @protocol microservice and to communicated encrypted traffic. A user microservice **124** or **146**, with its corresponding user container **122** or **148**, respectively, can be uniquely addressable using an assigned IP address and port number, as configured for the associated @sign in a corresponding root server entry.

(129) Container and microservice management, performed by the protocol provider server **103** and/or the registrar server **110**, can include management and deployment of user containers, in a multi-tenant container orchestration system. For example, user container management can include starting, stopping, and restoring user containers, deploying a given user microservice **124** or **146** as part of a user activation/registration process, and monitoring health of user containers, including resource allocation and usage, with automatic resource scaling as necessary. User containers can be reconfigured and redeployed in the event of an outage or hardware failure, for example. Data stored in user containers **122** or **148** can be encrypted for security. Decryption keys are discussed in more detail below. Users who choose to use a secondary server implemented as a user microservice **124** or **146** can pay fee(s) (e.g., activation, subscription fees) to the protocol provider or the registrar, respectively, as part of a registration process.

(130) Registration can also include a bootstrapping process in which a secondary server instance is

cryptographically bound to a user device **107**, such as during activation of the user device **107**, using a shared secret. For example, the registration application **138** can present a code (e.g., a QR (Quick Response) code) that includes the shared secret. The user can use the registration application **138** (or a camera application or similar functionality) to scan the code or otherwise retrieve the secret key from the code. The secret key can be stored on the user device **107** in a keys area **150**, preferably in a hardware trusted root on the user device **107**.

(131) The shared secret can also be provided (e.g., by the registrar server **110**) to the secondary server (e.g., microservice) associated with the user's @sign, for storage in the user's personal directory managed by the secondary server. The shared secret can be used by the user device **107** to authenticate to the secondary server. For example, an authentication process can include a key challenge and a challenge response, using the shared secret (e.g., using a CRAM (Challenge Response Authentication Mechanism) process).

(132) In some implementations, the shared secret is used to establish an authenticated connection between the user device **107** and the secondary server during activation and bootstrapping, with activation also including configuration of a public/private key pairing between the user device **107** and the secondary server. The user device **107** can generate a private key (and store the private key in the keys area **150**) and a corresponding public key, for sharing with the secondary server, such as by using a RSA (Rivest-Shamir-Adleman) algorithm. Once a public/private key pairing has been established, the shared secret can be removed from the secondary server (and in some cases from the keys area **150**). Removal of the shared secret from the secondary server can increase protection of encrypted user data on the secondary server, such as in a scenario where unauthorized access to the secondary server is somehow achieved. The public/private key pairing can be used for post-activation authentications of the user device **107** to the secondary server.

(133) Once the user device **107** is activated and configured, the user device **107** is recorded (e.g., in the @sign data **132**) as a registered device for updating data for the @sign handle to which the device **107** is registered. In some implementations, a given user device **107** can be registered as bound to multiple @signs. The user device **107** can store @sign handle(s) **152** for which the user device **107** is registered and corresponding public/private key pairs for each @sign.

(134) The settings application **145** running on the user device **107** can act as a secondary server for a registered @sign handle, receiving and responding to protocol requests locally on the user device. The settings application **145** can provide access to data stored in a personal directory **154** on the user device **107**, for example. When the user device **107** is bound to a secondary server, the settings application **145** can synchronize the personal directory **154** with a personal directory stored in the secondary server (e.g., in a user container **122** or **148**). When data in the personal directory **154** (or the corresponding personal directory in the secondary server) is updated, a “sync” protocol verb can be automatically invoked that updates the corresponding personal directory, so that user device-based and cloud-based directories are kept synchronized. The user device **107** can be authenticated to the secondary server before a synchronization occurs. The user device **107** can cache secondary server connection information in cached data **156** and/or the user device **107** can obtain secondary server connection information from the root server **105**.

(135) Synchronization can include one or more of sending data from one directory to the corresponding directory, comparing received items to stored items, identifying items in the receiver's directory that may need updating, and identifying changes that may have occurred in the receiver's directory that may in turn need updating in the sender's directory. Comparison processing can use metadata that indicates changes, times of changes, etc.

(136) Event-based synchronization for a network-connected user device **107** can ensure that changes occurring in one directory copy are reflected in the corresponding directory copy. If the user device **107** is occasionally offline, the user (and locally-installed applications) can still access the personal directory **154** on the user device **107** and any differences between the personal directory **154** and a secondary server directory copy can be automatically synchronized when the

user device **107** comes back online.

(137) If the user device **107** is offline, protocol exchanges can be routed to and handled by the secondary server. In some cases, for protocol exchanges that may require user permission, a notification/permission request can be sent by the secondary server to the user device **107**. If the user device **107** is offline, the permission request can be queued and a corresponding application/protocol request can be held until the user device **107** is back online and the user is able to respond to the permission request. In some implementations, the secondary server can be configured to hold a limited number (e.g., one) of permission requests/notifications for the user/user device **107**, so as to prevent a bad actor from overflowing the secondary server (or the user device **107**) with permission requests or other notifications. As another example, the secondary server can utilize a block list to prevent spamming. Such block lists may be device- or user-specific, or may be used across the entire protocol.

(138) The settings application **145** running on the user device **107** can enable a user of the user device **107** to configure access rights to data stored in the user's personal directory (e.g., in the personal directory **154** and in the mirrored secondary server copy). The user can make certain data public, private, or accessible only to certain users, or any other suitable limitations or restrictions. Other permissions applied to data can include whether data can be copied, shared, or used to derive other data. In general, access rights of various kinds can be granted or revoked using the settings application **145**. Access rights can be stored in association with corresponding personal data, such as in the personal directory **154** (and synced with the secondary server copy).

(139) As mentioned, the settings application **145** can serve protocol requests for user data, either from other protocol servers or from an application on the user device **107**. In general, the settings application **145**, along with the corresponding secondary server, can provide @sign server functionality for the user to enable the user device **107**, and the user's @sign, as protocol participants. The settings application **145** can present notifications to the user and enable the user to respond to notifications. Notification responses or other protocol responses can be provided to a sending protocol-enabled participant. Notifications and responses can be sent and received in a P2P fashion, and thus can be free from surveillance by a third party. Notifications can be received by the setting application **145** or the corresponding secondary server automatically, without a need for polling as used in other protocols such as POP3 (Post Office Protocol 3) or IMAP (Internet Message Access Protocol).

(140) In general, an @sign (e.g., protocol) server, whether running on a cloud server, dedicated hardware, or on a user device **107**, can provide an addressable endpoint for protocol-based communication. Along with a first, centralized tier of root servers **105**, protocol servers can provide a second, decentralized tier of the protocol architecture. While the root servers **105** can be configured to respond to simple @sign lookup requests, the more complex handling of protocol requests for data can be distributed to billions of separate decentralized resource directories (e.g., the protocol servers) that are responsible for responding to the requests. In other words, the protocol servers represent a fully decentralized set of personal service directories, with each protocol server configured to respond to lookup and other requests for a particular @sign. The response for a requested data item or service can come from a secure key/value persistence in the respective decentralized resource directory managed by the protocol server receiving the request.

(141) The protocol request received at a protocol server can include a general syntax pattern of "service@entity" that includes both the data or service requested being included in addition to the entity identifier, with the request being directed to a decentralized resource directory instance for processing. The response data for a particular service@entity request can be stored in a secure key/value persistence for the receiving entity. The value stored for the request can either be stored by value or by reference. The service@entity protocol syntax can be a base for creating Uniform Resource Identifiers (URIs). The service@entity syntax, while forming a simple base pattern, can be extended, such as to include an identity of a requester (e.g., @bob:phone@alice, which can be

provided to @bob by @alice in response to a request by @bob for information @alice has shared with @bob).

(142) The decentralized nature of the protocol servers can further increase security. Sensitive data can be stored in, for example, billions of separate, secure decentralized resource directory microservices. Data within a given microservice can be securely stored in an embedded key/value persistence with values strongly encrypted. Accordingly, sensitive data across users is not centralized or stored in one place, which can be security weaknesses for other types of systems. Decentralized distribution of data can support system scaling and can also reduce an attack surface for bad actors to a minimal footprint, as compared to other approaches.

(143) Storing data by reference can further increase data distribution and decentralization. With data being able to be stored either by value or by reference, a user's owned data can be distributed, using by-reference pointers, among various other systems, with access control (including encryption keys) still being stored and managed by the user's protocol server, which remains under the control of the user.

(144) Other than decentralization, various other approaches can be used for securing protocol requests and secured data when implementing the @protocol. For example, secure transport channels, such as TLS over TCP/IP, can be used. As another example, connections can be validated, by servers and/or clients, using SSL certificates. For further security, different protocol messages can be hashed and/or encrypted. Personal data can be encrypted and/or signed cryptographically, and can be decrypted by requesting and obtaining, for example, a public key of the entity. Other examples include use of a blacklist. A blacklist is a list of @signs that is checked when a "from:" verb is received by a protocol server. If the sender of the from: verb is found in the blacklist, the connection is immediately dropped, for example to fulfill a protocol promise of preventing a known blocked @sign from having any interaction whatsoever with the protocol server apart from a public lookup (e.g., which does not use the from: verb). In addition to a blacklist, "allowed list(s)" or "whitelist(s)" can be maintained by a protocol server that are derived from shared data and include @signs to whom the shared data is shared.

(145) Additionally, the protocol servers are designed to prevent outside surveillance. Data stored in a protocol server can only be accessed by the registrant and other users to which the registrant user has explicitly shared access. The protocol servers are specifically designed to control the granting or revoking of access rights to registrant's data and services in real-time, based on the user's permissions. Other tracking or access to a protocol server outside the protocol specification can be prohibited, even from the registrar server **110**, the protocol provider server **103**, or other components of the system **100**.

(146) The personalized decentralized resource directories provided by the various protocol servers can enable various types of applications, including P2P applications. For example, a user device **107** can include various protocol-enabled applications **158**. The protocol-enabled applications **158**, which can include installed and/or web-based applications, can retrieve an @handle from the settings application **145**, request approval for access to data in the personal directory **154**, and expose application data **160** to the settings application **145** (or the secondary server) for the user's protocol server to maintain permissions for the application data **160**. Although shown separately from the personal directory **154**, application data **160** can be stored in the personal directory **154** or linked from the personal directory **154**.

(147) In some implementations, a protocol-enabled application **158** (and/or the settings application **145**) includes a resolver library which can make protocol requests on behalf of the protocol-enabled application **158**. Protocol-enabled applications **158**, through use of the resolver library and/or the settings application **145**, can be configured to receive user input, translate user input into protocol commands, process protocol verb requests, provide protocol verb responses, and persist data. Data persisted by an application can be data shared by others, and can be stored in the application data **160**, the personal directory **154** and/or the cached data **156**. Data shared by others can have a TTL

(Time to Live) setting, which limits how long the data from others may be accessible.

(148) Data in the personal directory **154**, the personal directory **144**, or in secondary server personal directory copies can also have a TTL setting. As described below, a private lookup challenge value can be placed in a personalized directory with a TTL value. As another example, application and/or personalized directory data can have a “TTB” (e.g., Time unTil Birth, Time To Birth) setting, which can represent a duration value. Data with a TTB setting can be configured to become available in the application data or personalized directory after the duration has elapsed (e.g., sharing information with a TTB parameter can provide access to the information only after the TTB time has been reached). For example, important financial information might be stored in an @server that is to become accessible only at a defined time in the future for all recipients. The financial information can be public or private and might also be subject to other time based controls, such as TTL and TTR (Time To Refresh). For instance, data can have both a TTB and a TTL setting, which can result in a data item appearing after a TTB duration has expired, and disappearing after a TTL duration has expired. Data with a TTR setting can be refreshed when the interval value for the TTR setting is reached. TTL, TTB, and TTR settings, whether used separately or in combination(s), can enable flexible and powerful data access and privacy control capabilities.

(149) Applications can present data in a GUI (Graphical User Interface) **162**. As another example, the protocol-enabled applications **158** can include a shell application, in which a user can enter protocol commands, e.g., after connecting to a root server **105**, or a protocol server. The user can enter anonymous commands to access publicly available information, and can authenticate to a given protocol server, to make an authenticated request for information.

(150) A protocol-enabled application **158** can be a standalone application for managing personal data, or can be a P2P application that communicates with other P2P applications, over the network **101**. End-to-end encrypted P2P applications can have an advantage of not being subject to surveillance as they do not need to communicate through backend servers, or send unencrypted information in the clear. Rather, P2P applications can communicate in a distributed fashion, directly between themselves, at the edge of the network **101**.

(151) The protocol provider can support the application developer community, by providing the source code distribution **114** and the documentation distribution **116**. The source code distribution **114** can be an open source code repository, e.g., for reference implementations, and can be made available from the protocol provider server **103** (or another server of the protocol provider) or through another source code repository or channel. The source code distribution **114** can include code for a SDK (Software Development Kit) that can be used by application developers **104** to create protocol-enabled applications. Other source code examples include sample code for protocol exchanges, code for fully-developed sample applications, etc.

(152) The documentation distribution **116** can include protocol specification information and documentation, discussion of example protocol exchanges, tutorials, or other types of documentation. In addition to source code and documentation, the protocol provider can provide binary distributions for development tools, sample applications, etc. In general, an overall repository managed or supported by the protocol provider can include items that can be used, tested, evaluated, contributed to, and extended, by the application developer community.

(153) Application developers **104** can use developer tools **164** provided by the protocol provider (or by other providers) to develop protocol-enabled applications. Application developers **104**, in order to use the protocol, can first agree to terms of conduct published by the protocol provider, which can provide legal recourse. Application developers can review application certification guidelines included in the documentation distribution **116**, to learn requirements for application certification. When an application is ready for certification, an application developer **104** can provide application code **166** (or another type of packaged application distribution) to the application certifier **118**, for certification.

(154) Application certification can include review and testing (e.g., by the application certifier **118**

or by administrator(s)) of the application code **166**, to check for compliance with protocol guidelines. Upon successful certification of an application, application certification information **119** can be stored for the certified application. As another example, certification information can be stored at an @sign server associated with the application. In some implementations, application certification can be validated in real time, such as by using a checksum validation. Upon receiving an application certification result, an application developer **104** (or the application certifier **118**) may update code of the application so that, for example, a user interface of the application presents an indication of certification to a user.

(155) As another example, application information stored for the application in an application repository **168** of the application store **108** can be updated so that prospective users can see the certification status of the application. A user device **107** can connect to an application store service **170** (e.g., using a web site or an installed application provided by the application store **108**), to access various application discovery tools. For example, a user can use the application store service **170** to browse and/or search for protocol-enabled applications. As another example, the application store service **170** can provide application ratings and/or recommendations. Once a user selects an application, the application can be downloaded (e.g., to a user device **107**) from the application repository **168**. The application store **108** can be provided by or otherwise affiliated with the protocol provider. As another example, at least some other application stores can be provided by third party vendors.

(156) In some implementations, the backup server **106** is included in the system **100**, as a separate backup repository for protocol servers. For example, a backup service **171** can provide backup/restore access to personal directory backups **172**. The personal directory backups **172** can hold non-key data, for example. In some implementations, backup services can be provided by secondary server providers (e.g., data may be backed up in the user containers **122** or **148**). If a user chooses to use the backup server **106** or another backup mechanism, the user can be charged a fee for the backup services.

(157) In some implementations, the enterprise user device **111** can be connected over the local area network **112** with the enterprise microservice server **113** that can act as a proxy for the enterprise user device **111** for network connectivity beyond the local area network **112**. The enterprise user device **111** may be behind a corporate or other type of firewall, for example. The enterprise user device **111** can send a “plookup” (e.g., a proxied lookup) protocol verb to the enterprise microservice server **113**, for protocol exchanges. Through use of the plookup verb, the enterprise microservice server **113** can contact other protocol servers for protocol exchanges to be performed on behalf of the enterprise user device **111**. The enterprise microservice server **113** can be useful in large enterprise environments that prefer all lookups going through a single Internet-connected secondary server, or where only a single port (e.g., on the enterprise microservice server **113**) is preferred to be opened through the firewall for protocol use, or for other cases where a minimizing of connections is desired. The enterprise microservice server **113** can support an internal, private @protocol system (that includes the enterprise user device **111** and other devices), that is located internally within the enterprise.

(158) As used in the present disclosure, the term “computer” is intended to encompass any suitable processing device. Indeed, the illustrated devices and systems may be or include any computer or processing device such as, for example, a mobile device, a laptop, a blade server, a general-purpose personal computer (PC), Mac®, workstation, UNIX-based workstation, IoT device, or any other suitable device. In other words, the present disclosure contemplates computers other than general purpose computers, as well as computers without conventional operating systems. Further, the illustrated devices and systems may be adapted to execute any operating system, including Linux, UNIX, Windows, Mac OS®, Java™, Android™, iOS, Dart, Swift, Kotlin, or any other suitable operating system.

(159) Interfaces **175**, **176**, **177**, **178**, **179**, **180**, **181**, **182**, **183** and **184** (referred to below as

interfaces **175-184**) are used by the governing body **102**, the protocol provider server **103**, the application developers **104**, the root servers **105**, the backup servers **106**, the user devices **107**, the application store **108**, the protocol servers **109**, the registrar servers **110**, and the enterprise microservice server **113**, respectively, for communicating with other systems in a distributed environment—including within the system **100**—connected to the network **101**. Additionally, the interface **185** enables the enterprise user device **111** to communicate on the LAN **112** (e.g., to other enterprise user devices **111** and/or to the enterprise microservice server **113**). Generally, the interfaces **175-184** each comprise logic encoded in software and/or hardware in a suitable combination and operable to communicate with the network **101**. More specifically, the interfaces **175-184** may each comprise software supporting one or more communication protocols associated with communications such that the network **101** or interface's hardware is operable to communicate physical signals within and outside of the illustrated system **100**.

(160) The protocol provider server **103**, the root servers **105**, the backup servers **106**, the user devices **107**, the protocol servers **109**, and the registrar servers **110** may each respectively include one or more processors **186, 187, 188, 189, 190, or 191**. Each processor **186, 187, 188, 189, 190, or 191** may be a central processing unit (CPU), a blade, an application specific integrated circuit (ASIC), a field-programmable gate array (FPGA), or another suitable component. Generally, each processor **186, 187, 188, 189, 190, or 191** executes instructions and manipulates data to perform the operations of the respective device. Specifically, each processor **186, 187, 188, 189, 190, or 191** executes the functionality required to receive and respond to requests from another device.

Although not illustrated as including processors, the governing body **102**, the application developer **104**, the application store **108**, the enterprise user device **111**, and the enterprise microservice server **113** may each include (or be associated with device(s) that include) one or more processors.

(161) Regardless of the particular implementation, “software” may include computer-readable instructions, firmware, wired and/or programmed hardware, or any combination thereof on a tangible medium (transitory or non-transitory, as appropriate) operable when executed to perform at least the processes and operations described herein. Indeed, each software component may be fully or partially written or described in any appropriate computer language including C, C++, Java™, JavaScript®, Visual Basic, assembler, Peri®, Dart, Swift, Kotlin, Go, Python, any suitable version of 4GL, as well as others. While portions of the software illustrated in FIG. **1** are shown as individual modules that implement the various features and functionality through various objects, methods, or other processes, the software may instead include a number of sub-modules, third-party services, components, libraries, and such, as appropriate. Conversely, the features and functionality of various components can be combined into single components as appropriate.

(162) Each of the devices or systems illustrated in FIG. **1** may include one or more memories (e.g., storage locations). Each memory may include any type of memory or database module and may take the form of volatile and/or non-volatile memory including, without limitation, magnetic media, optical media, random access memory (RAM), read-only memory (ROM), removable media, or any other suitable local or remote memory component. Each memory may store various objects or data, including caches, classes, frameworks, applications, backup data, business objects, jobs, web pages, web page templates, database tables, database queries, repositories storing business and/or dynamic information, and any other appropriate information including any parameters, variables, algorithms, instructions, rules, constraints, or references thereto associated with the purposes of the system **100**.

(163) The user devices **107** may each generally be any computing device operable to connect to or communicate with other devices using a wireline or wireless connection, with different user devices **107** being similar or different types of devices in a particular implementation. In general, the user device **107** comprises an electronic computer device operable to receive, transmit, process, and store any appropriate data associated with the system **100** of FIG. **1**. The user device **107** can each include one or more client applications, including the registration application **138**, the settings

application **145**, and the protocol-enabled applications **158**. A client application is any type of application that allows the user device **107** to request and view content on user device **107**. In some instances, a client application may be an agent or client-side version of the one or more server applications running on a server.

(164) The user device **107** is generally intended to encompass any client computing device such as a laptop/notebook computer, wireless data port, smart phone, personal data assistant (PDA), tablet computing device, one or more processors within these devices, or any other suitable processing device. For example, the user device **107** may comprise a computer that includes an input device, such as a keypad, touch screen, or other device that can accept user information, and an output device that conveys information associated with the operation of the client device itself or the system **100**, including digital data, visual information, or the GUI **162**.

(165) The GUI **162** interfaces with at least a portion of the system **100** for any suitable purpose, including generating a visual representation of the registration application **138**, the settings application **145**, or the protocol-enabled applications **158**. In particular, the GUI **162** may be used to view and navigate various Web pages. Generally, the GUI **162** provides a respective user with an efficient and user-friendly presentation of data provided by or communicated within the system **100**. The GUI **162** may comprise a plurality of customizable frames or views having interactive fields, pull-down lists, and buttons operated by the user. The GUI **162** contemplates any suitable graphical user interface, such as a combination of a generic web browser, intelligent engine, and command line interface (CLI) that processes information and efficiently presents the results to the user visually.

(166) There may be any number of user devices **107** associated with, or external to, the system **100**. Additionally, there may also be one or more additional user devices **107** external to the illustrated portion of system **100** that are capable of interacting with the system **100** via the network **101**. Further, the term “client”, “client device” and “user” may be used interchangeably as appropriate without departing from the scope of this disclosure.

(167) FIG. 2A is a block diagram illustrating an example system **200** for user-controlled access control for user information. The system **200** may be similar to at least portions of the system **100**, for example. The system **200** includes a user device **202** of a “Bob” user and a user device **204** of an “Alice” user. Various figures below describe various scenarios for data sharing and access between or by an example “Bob” and/or “Alice” user. As described below with respect to FIGS. 14-20C, the user devices **202** and **204** can be onboarded to the @sign platform. The @sign platform, as described above, can include a namespace directory and decentralized resource directory servers that communicate over a network **205** (e.g., the Internet).

(168) The namespace directory can be implemented as a root server **206**. Decentralized resource directory servers can be implemented as secondary servers, including a secondary server **207** for the user “Bob” and a secondary server **208** for the user “Alice.” The root server **206** can provide a lookup service **210** that accepts lookup requests for secondary server information for a particular @sign. For example, the root server **206** can maintain, (e.g., in an in-memory database **211** (or other type of database), entries **212** that are in a “domain:port” format **214**. For example, the entries **212** include an entry **216** for the secondary server **207** for the user “Bob” and an entry **218** for the secondary server **208** for the user “Alice.”

(169) The secondary servers **207** and **208** can each manage a decentralized resource directory for a respective user, such as using a secure key/value persistence. Each secondary server **207** and **208** can provide responses to requests for user data managed by the respective secondary server. The secondary servers **207** and **208** can correspond, for example, to the @sign server **109** of FIG. 1. For example, a same type of server can be referred to using different names in different implementations. For instance, @sign server, secondary server, protocol server, DESS (Distributed Edge Secondary Server), microserver, or other names can be given to a server that services an individual decentralized resource directory. As described in more detail below, individual

decentralized resource directories can also be implemented locally on the user device **202** or **204**, with synchronization performed as needed between the secondary server **207** or **208**, respectively. For example, the secondary servers **207** and **208** can be cloud-based secondary servers and can communicate with the user device **202** or the user device **204** over a cellular (or other) network **220** or **222**, respectively. In some implementations, a cloud secondary server such as the secondary server **207** and **208** can be referred to using the term “secondary” in that a secondary server manages a cloud-based individual decentralized resource directory that is secondary to (e.g., a backup to) a local individual decentralized resource directory that is local to a user client device.

(170) As part of onboarding a respective user device **202** or **204**, a shared secret shared between the user device **202** or **204** and the corresponding secondary server **207** or **208** can be provided (e.g., in various ways, as described below) to the user device **202** or **204** by a registrar **224**. The shared secret can be used to establish a secured relationship between the user device **202** or **204** and the corresponding respective secondary server **207** or **208**. Once a relationship has been established between the user device **202** or **204** and the corresponding respective secondary server **207** or **208**, cryptographic key generation **226** or **228** can be performed on the user device **202** or **204**, respectively, and use of cryptographic keys can replace use of the shared secret. The cryptographic key generation **226** and **228** can include generation of a public/private key pair. A generated private key can be stored on the user device **202** or **204** (e.g., in a hardware entrusted root). A generated public key can be shared by the user device **202** or **204** with the respective corresponding secondary server **207** or **208**. Data sharing between the user device **202** or **204** and the corresponding respective secondary server **207** or **208** can be performed using the @protocol, as described below. Additionally, once configured, the secondary servers **207** and **208** can perform data sharing and data requests using the @protocol, over the network **205**, with any protocol-enabled device. For example, private and public lookups and data sharing can be performed.

(171) FIG. 2B is a block diagram illustrating an example system **250** for a public lookup transaction. A public lookup can be performed on an @sign server without being authenticated to the @sign server. In this example, a user **251** (“Bob”) wants to do a lookup for a location value for an @alice user. The user **251** can perform the lookup using a device **252** (e.g., Bob's device, configured for or associated with an @bob user). The device **252** can include code for executing @protocol exchanges (e.g., as provided by a resolver library), among other protocol configurations.

(172) In a first stage (e.g., illustrated as a circled “one”), the device **252** sends an “@alice” query **254** to a root server **256**. A public lookup of an @sign can be performed without being authenticated to the root server **256**. The root server **256** looks up server information for the @alice handle and, in stage two, returns server address information **258** to the device **252**. The server address information **258** can include a server address and port number for a protocol server **260** that handles protocol requests for the @alice user.

(173) In a third stage, the device **252** uses the server address information **258** to send a “location@alice” query **262** to the protocol server **260**. The protocol server **260** can determine whether a “location” value is stored by the protocol server **260** and whether the location value is publicly accessible. The query **262** is an anonymous, unauthenticated (e.g., public) query, so the protocol server **260** can be configured to provide a response to the query **262** only if the protocol server **260** stores a public “location” value. In some implementations, if the protocol server **260** does not store a public location value, the protocol server **260** can provide a null (e.g., empty) response. Alternatively, if no public location value is stored, then the protocol server **260** can ignore the location@alice query **262** with no response provided.

(174) In a fourth stage, in response to determining that a public location value is stored at the protocol server **260**, the protocol server **260** sends a retrieved public location value of “California” to the device **252**. A note **266** summarizes the transaction (i.e., the device **252** receiving the public location value of “California” **264** in response to the public-lookup “location@alice” query **262**).

(175) The lookup:location@alice query **262** is an example of a lookup protocol verb used to lookup

a particular address in a particular @handle's namespace. The protocol server **260** can be configured to assume that any received value that ends with “@alice” is a lookup request, for example. As mentioned, an unauthenticated request is referred to as a public lookup, in which a requesting user does not have to prove who they are. Private, authenticated lookups are also supported, as described below.

(176) FIG. **3** is a block diagram illustrating an example system **300** for a private lookup transaction. The @protocol includes an automated method for ensuring that an entity that is requesting some private information from another entity has proven that they are who they claim to be before the request is presented for processing. Failure of this proof can immediately cause the connection between the entities to be dropped. A “proof of life” or “pol” algorithm can be performed during a private lookup process.

(177) A device **301**, associated here with user @bob, can include a resolver **302**. The resolver **302** can be a library included in (or otherwise used by) a protocol-enabled application. The resolver **302** can be configured to perform (among other tasks) a private lookup of information managed by a user with a handle of @alice. The resolver **302** can be included in an application that is running on the device **301** on behalf of the user @bob.

(178) In a first stage, the resolver **302** sends an “@alice” query **304** and an “@bob” query **306** to a root server **308**. The resolver **302** may or may not currently have access to server information for a protocol server for @bob, for example. To perform the pol algorithm as part of a private lookup for information managed by the @alice user, the resolver **302** can connect to protocol servers for both @alice and @bob.

(179) In response to the @alice query **304** and the @bob query **306**, the root server **308** looks up server information for the @alice and @bob handles and, in a second stage, returns, to the resolver **302**, server address information **310** for a protocol server **311** for @alice and server address information **312** for a protocol server **314** for @bob, respectively.

(180) In a third stage, the resolver **302** submits a “from:@bob” request **316** to the protocol server **311**. The “from:@bob” request **316** is an example use of a from protocol verb that can be used to tell a protocol server what @handle is associated with a request. The pol algorithm can be initiated when a user submits a from verb that includes an @sign of the requester. The “from:@bob” request **316** includes the @sign “@bob” in this example.

(181) In the pol algorithm, upon receiving a “from” request, the requested entity first returns a challenge asking the requestor to publish a defined secret in a defined secret location in the requestor's namespace and notify the requested entity when the defined secret is available at the defined location. For example, in a fourth stage, the protocol server **311**, in response to the request **316**, sends a challenge requesting the resolver **302** to place a cookie **318** at a particular address **320** on the protocol server **314**. That is, proof of a requester's identity can be provided using an ability to place a cookie or other identifier at a specified place in the requester's @handle name space. For the resolver **302** to legitimately act on behalf of @bob, the resolver **302** should be able to authenticate to the protocol server **314** for @bob and place the requested cookie **318** at the requested address **320**. The cookie **318** and the address **320** can be randomly-generated values generated by the protocol server **311** as part of the pol algorithm. The cookie **318** and the address **320** can be UUIDs (Universally Unique IDentifiers) or some other type of random value or identifier.

(182) In a fifth stage, the resolver **302** authenticates to the protocol server **314** for @bob and places a cookie **322** (corresponding to the requested cookie **318**) at an address **324** (corresponding to the requested address **320**). The resolver **302** authenticates to the protocol server **314** before performing the placement of the cookie **322**. Different approaches can be used for authenticating to the protocol server **314**.

(183) For example, the resolver **302** can first indicate to the protocol server **314** that the resolver **302** is acting on behalf of the @bob user (e.g., using a “from:@bob” command). The protocol

server **314** can respond with a challenge for the resolver **302**. The resolver **302** can respond to the challenge, using for example, a CRAM verb or a PKAM (Public Key Authentication Mechanism) verb. The CRAM verb uses a shared secret, while the PKAM verb uses a public/private key approach. Authentication is discussed in more detail below with respect to FIG. 4A.

(184) After successfully authenticating to the protocol server **314**, the resolver **302** can submit an update request to the protocol server **314** to place the cookie **322** at the address **324**. An update verb can be used to place a new value or update an existing value on the protocol server **314**, for example. Placing the cookie **322** at the address **324** can prove to the protocol server **311** that the resolver **302** is in fact acting on behalf of the @bob user.

(185) In a sixth stage, the resolver **302** sends an indication to the protocol server **311** that the requested cookie is in place on the protocol server **314** and available for testing by the protocol server **311**. For example, the resolver **302** can send a pol verb to the protocol server **311**. The pol verb can be used to signal the protocol server **311** to check for (e.g., look up) the cookie **318** on the protocol server **314**.

(186) In a seventh stage, the protocol server **311** sends an “@bob” query **325** to the root server **308** (or to another root server) requesting protocol server information for the @bob handle.

(187) In response to the @bob query **325**, the root server **308** looks up server information for the @bob handle and, in an eighth stage, returns, to the protocol server **311**, server address information **326** for the protocol server **314** for @bob.

(188) In a ninth stage, the protocol server **311** looks for the cookie **318** at the address **320** on the protocol server **314**. Looking for the cookie can include sending a lookup request for a value on the protocol server **314** at the address **320**.

(189) In a tenth stage, the protocol server **314** provides the cookie value **322** stored at the address **324** to the protocol server **311**, in response to the lookup request sent by the protocol server **311**.

(190) In an eleventh stage, the protocol server **311** “tastes” the cookie received from the protocol server **314**. That is, the protocol server **311** compares the cookie **318** sent in the proof challenge to the cookie **322** that was actually stored at the address **320** on the protocol server **314** (and sent to the protocol server **311** by the protocol server **314**). If the cookie **318** does not match the cookie **322**, the protocol server **311** can determine that the resolver **302** has failed the proof of life challenge. The protocol server **311** can drop a connection with the resolver **302**, simply not respond to the pol verb sent by the resolver **302** in the sixth stage, or return an error code to the resolver **302** that indicates a failure of the pol attempt. Similarly, the protocol server **311** can drop a connection, ignore the pol verb, or send an error code to the resolver **302** if the cookie **318** was simply not found on the protocol server **314**.

(191) In a twelfth stage, in response to a successful retrieval and test of the cookie **322**, the protocol server **311** sends an indication of pol success to the resolver **302**, indicating that the resolver **302** has successfully authenticated to the protocol server **311** as the user @bob.

(192) In a thirteenth stage, the resolver **302** sends, using the authenticated connection, a “lookup location@alice” request **327** to the protocol server **311**. The protocol server **311** can determine that a value specific to @bob is stored by the protocol server **311** for the location@alice address.

(193) In a fourteenth stage, the protocol server **311** returns a value of “Home” **328** to the resolver **302**, in response to the lookup request **327**. A note **330** summarizes the private lookup transaction, of the resolver **302** receiving the location value of “Home” **328**, in response to the private (e.g., as @bob) lookup “location@alice” request **327**. That is, the value of “Home” has been shared with the @bob user and is provided in response to location@alice requests sent by the @bob user or an application acting on behalf of @bob, in contrast to a publicly accessible location value **214** of “California” that is returned in response to location@alice requests sent by unauthorized users (or by users to which a private/custom location value is not shared).

(194) FIGS. 4A and 4B are a combined swim lane diagram illustrating an example method **400** for a private lookup transaction. Similar to the private lookup illustrated in FIG. 3, the method **400** can

describe a user “Bob” with a handle of “@bob” performing a private lookup of information maintained by a user “Alice” with a handle of “alice.”

(195) As shown in FIG. 4A at **402**, a resolver **404** sends a query **406** of “@alice” to a root server **408**. At **410**, the root server **408** sends a response **412** to the query **406** of “alice.alserver.com:6455” that includes a server and port number of Alice's server **413**. Bob can use an application that includes, implements, or accesses the resolver **404** to submit a query for Alice's server information, if Bob (or the application) does not have Alice's server information cached. If Bob (or Bob's application) does not have his own server information cached, the resolver **404** can, at **414**, send a query **416** of “@bob” to the root server **408** to obtain server information for Bob's server **417**. For example, at **418**, the root server **408** sends a response **420** of “bob.bobserver.com:6200” in response to the query **414**.

(196) In general, the root server **408** can respond to a query for an @sign with a single IP address and port or a single fully qualified domain name (FQDN) and port number. The returned port number can identify the @protocol service and a particular user's server. The root server **408** can respond with a null value (or simply ignore a request) if the root server **408** does not have server information for a particular requested @sign.

(197) At **422**, the resolver **404** connects to Alice's server **413** and sends a “from:@bob” message **426** to Alice's server **413**. As discussed above, a “from” request can indicate to a requested entity what @handle the requester claims to be. The message **426** can indicate to Alice's server **413** that the resolver **404** is claiming to be associated with the @bob user.

(198) At **427**, in response to the message **426**, Alice's server **413** generates pol challenge information. For example, Alice's server **413** can generate a first random value and a second random value to be used in a challenge sent to the resolver **404**.

(199) At **428**, Alice's server **413** responds to the message **426** by sending a pol challenge **430**. The pol challenge **430** is a proof challenge that requests the resolver **404** to place a first identifier (e.g., a UUID) “271b7d94-8a31-48e3-9348-daf909349b21” at a randomly-generated address of “86adcd50-9fcb-4e76-a17f-1a572b6a1ee5” on Bob's server **417**. If the resolver **404** is authorized to act on behalf of @bob, the resolver **404** should be able to successfully perform the requested challenge. Although the illustrated pol challenge **430** uses a “proof” keyword (e.g., proof verb), other keywords can be used. For instance, in some implementations, a “data” keyword is used rather than a “proof” keyword. Likewise, for other protocol verbs or syntax, different variations can be used, in different implementations.

(200) At **432**, the resolver **404** sends a “from:@bob” message **434** to Bob's server **417** to attempt to authenticate to Bob's server **417**. Similar to the message **426**, the message **434** is indicating to Bob's server **417** that the resolver **404** is claiming to be acting on behalf of the @bob user.

(201) At **437**, Bob's server **417** generates key challenge information.

(202) At **438**, Bob's server **417** sends a key challenge **440** to the resolver **404**. As shown, if a user, application, or process issues a “from:@bob” verb to the server associated with the @bob user, the server for @bob can issue a key challenge, whereas the server for @bob can issue a proof challenge in response to a from: verb from a user other than the @bob user. As discussed above, the resolver **404** can respond to the key challenge **440** using a CRAM verb, a PKAM verb, or some other type of authentication response.

(203) At **442**, the resolver **404** (or another process, on behalf of the resolver **404**) executes a PKAM command **444** to create a signed version of the key challenge **440** using a private key for the @bob user. If the resolver **404** is legitimately configured to represent the @bob user, the resolver **404** has access to the private key for the @bob user. With the PKAM approach, Bob's server **417** is configured to have access to the public key for the @bob user.

(204) Referring now to FIG. 4B, at **446**, the resolver **404** sends a PKAM message **448** to Bob's server **417**. The PKAM message **448** includes a signed version of the key challenge **440** that was created by the PKAM command **444**.

(205) At **450**, Bob's server **417** tests the PKAM message **448**. For example, Bob's server **417** can check to see if the signed version of the key challenge **440** is accurate by testing the signed version of the key challenge **440** using the public key for the @bob user.

(206) At **452**, in response to a successful test of the PKAM message **448**, Bob's server **417** sends an indication **454** of authentication success to the resolver **404**.

(207) At **456**, the resolver **404** sends, while authenticated on Bob's server **417** as @bob, an update request **458** to Bob's server **417** to place the first random value "271b7d94-8a31-48e3-9348-daP909349b21" included in the pol challenge **430** at the randomly-generated address of "86adcd50-9fcb-4e76-al7f-1a572b6a1ee5 included in the pol challenge **430**.

(208) The update protocol verb enables a user or process authenticated on a given @sign server to update entries in that @sign namespace. A general format of the payload for an update verb is: [public/@sign]:key@[@yourname] value. An example update request is "update:public:email@alice alice.bloggs@emailXXX.com," which can create a public record for an email@alice address with a value of "alice.bloggs@emailXXX.com." An example of updating a value shared to a certain user is "update:@bob:email@alice alice@personalemail.org." A delete verb can be used by an authenticated user or process to delete an entry on an @sign server. For example, "delete:public:email@alice" and "delete:@bob:email@alice" requests can be sent to delete the entries mentioned in the previous two examples.

(209) Continuing with FIG. **4B**, at **460**, Bob's server **417** performs an update of Bob's personal directory based on the update request **458**. In some implementations, a cookie placed on Bob's server based on the update request **458** is configured with a timeout value. For example, the cookie can have a TTL (Time To Live) value set to one minute. After one minute has passed, Bob's server **417** can delete the cookie (or mark the cookie as invalid). A timeout value, combined with generation of random values such as UUIDs, can result in a negligible likelihood of a namespace clash (e.g., of another application or user placing a same cookie value at a same location). From a perspective of other entities other than the resolver **404** and Bob's server **417**, the update request **458** essentially results in the placing of an unknowable value at an unknowable location. A likelihood of another user or application knowing (or learning) the random cookie value or the random location can be substantially negligible. In some implementations, Bob's server cryptographically signs the cookie using a private key associated with the @bob user before placing an encrypted cookie at the address requested by Alice's server **413**. Alice's server **413** can decrypt the signed cookie, as described below.

(210) At **462**, Bob's server **417** sends an indication **464** of a successful update to the resolver **404**. At **466**, the resolver **404** sends a pol request **468** to Alice's server **413**. The pol request **468** can be used to signal Alice's server **413** that Bob's server **417** is ready for Alice's server **413** to verify the pol challenge **430**.

(211) At **470**, Alice's server **413** sends a query **472** of "@bob" to the root server **408** to obtain server information for Bob's server **417**. At **474**, the root server **408** sends a response **476** of "bob.bobserver.com:6200" in response to the query **472**.

(212) At **478**, Alice's server **413** sends a lookup request **480** to Bob's server **417** to look up the value at the address Alice's server included in the pol challenge **430**.

(213) At **482**, Bob's server **417** sends a value **484** to Alice's server **413** in response to the lookup request **480**. The value **484** is retrieved by Bob's server **417** as the value of the 86adcd50-9fcb-4e76-al7f-1a572b6a1ee5@bob address on Bob's server **417**.

(214) At **486**, Alice's server **413** performs a proof test. Alice's server **413** can compare the value **484** to the value Alice's server **413** included in the pol challenge **430**. In implementations in which Bob's server **417** cryptographically signs the value Alice's server **413** included in the pol challenge **430**, before storing a digitally-signed value on Bob's server **417**, Alice's server **413** can decrypt the signed value using a public key for the @bob user, before comparing a decrypted value to the value included in the pol challenge **430**. Alice's server **413** can obtain the public key for @bob by doing a

lookup for publickey@bob, for example.

(215) At **488**, in response to a successful proof test, Alice's server **413** sends an indication **490** of authentication success to the resolver **404**. The resolver **404** is now authenticated as the @bob user on Alice's server **413**. Alice's server **413** has concluded that the resolver **404** is in fact acting on behalf of the @bob user as claimed by the resolver in the "from:@bob" message **426**.

(216) At **492**, the resolver **404**, while in an authenticated state with respect to Alice's server **413**, sends an authenticated lookup request **494** to Alice's server **413**.

(217) At **496**, Alice's server **413** sends a user-specific value **498** to the resolver **404**, in response to the authenticated lookup request **494**. The user-specific value **498** is an example of a polymorphic response. Alice's server can send different values in response to a same lookup location@alice request, based on whether the sender is anonymous or based on the identity of the sender.

(218) FIG. 5 illustrates example keys **500** of values stored in a user's personal directory. The user's personal directory, for example, may be similar to the personal directory **144** and/or the personal directory **154** of FIG. 1. The keys **500** can be for a user Alice, with an address of @alice. A key **502** can be for an email password that Alice saves for herself only. Other keys can be for values that are shared with specific users. For example, a key **504** can be for a location value that is shared with an @barbara user, a key **506** can be for a phone number that is shared with @barbara, and a key **508** can be for a phone number that is shared with an @bob user. The phone number shared with @barbara can be a different phone number than the phone number shared with @bob. An item of data can be shared to multiple users (e.g., a group of users). A respective grant for the item of data can be granted to each user in the group, for example.

(219) Keys **510** and **512** are for private (e.g., internal) synchronization and private-key values, respectively. The keys **500** may be stored on Alice's device (e.g., a mobile device), for example, and the keys **510** and **512** may be relevant for protocol activity on the device. Keys **514** and **516** can be for a publicly published email address and location values, respectively.

(220) The key **518** can be for a public "home city" value. Rather than use the word "home" in the key **518**, the key **518** includes a symbolic house character. Keys (and values and @sign handles) can be a combination of characters from a Unicode character set, for example. For instance, Unicode can support 137,994 characters covering 150 modern and historic scripts (e.g., beyond the Latin script), at least **64** languages, as well as multiple symbol sets, including emoji and graphical characters. Accordingly, users have significant flexibility for being creative with selecting keys, values, and their @handle address.

(221) Some rules can be defined for legal @handle addresses, however. For instance, white space characters and certain predefined control characters and punctuation characters can be defined as illegal for an @handle address. Additionally, @handle addresses can be case insensitive, so that "@alice" and "@Alice" refer to the same @sign and same user.

(222) FIGS. 6A-6C illustrate example uses of a scan verb. The scan protocol verb can be used to scan available @-addresses (e.g., keys) for a user or application either at a public level or once a pol process has been completed. The scan verb can be used to discover and harvest available @-addresses. As shown by a message **602** in FIG. 6A, a connection has been established (by a particular user or application) to Alice's protocol server.

(223) A scan verb **604** is entered (or otherwise sent) to Alice's server. In response to the scan verb **604**, Alice's server determines available @-addresses for a current unauthenticated state. That is, given the unauthenticated state of the connection, Alice's server can identify and return, in a response **606**, publicly available @-addresses on Alice's server. For instance and as shown in the response **606**, publicly available @-addresses include "email@alice" and "location@alice."

(224) In some implementations, a server can exclude certain @-addresses from being included in a response to a scan verb. For example, addresses that begin with a special character (or characters), such as "_", can be excluded. Addresses that begin with "_" can be considered secret, in that they are not included in a scan response but can be looked up if known. For instance, a lookup request

608 for “secret@alice” can be responded to with an output **610**. A service address that starts with “_” can be publicly available for lookup only if the requesting entity is aware of the service address. The entity providing the service can selectively share the service address, for example. The service address can be or include a value, such as a long and/or random string of characters, which would be impractical for other entities to determine. For example, an @bob user can store a location value for a party at an example service address of _partyforbob_14445aa444fc@bob. Bob can share the address with those who are invited to the party. Other entities who are not aware of the service address would not know that a value is stored at that service address.

(225) As shown in FIG. 6B, a user Bob has entered a “from:@bob” command **622** while connected to Alice's server, and in response to a proof challenge **624**, has entered a pol command **626** (e.g., after placing a requested identifier at a requested location on Bob's server). An “@bob@” prompt **628** indicates that the pol command **626** was successful and that Bob is now authenticated as Bob on Alice's server. While authenticated as Bob, a scan command **630** can result in output **632** that lists @-addresses on Alice's server that are available to Bob, including “email@alice,” “location@alice,” and “phone@alice” addresses. The output **632** includes the public keys described above with respect to FIG. 6A as well as a “phone@alice” key of a phone number value that has been shared by Alice with Bob.

(226) As shown in FIG. 6C, Alice has entered a “from:@alice” command **642** while connected to Alice's server, and in response to a key challenge **644**, has entered a CRAM command **646** (e.g., after determining a correct digest to include with the CRAM command **646**). An “@alice@” prompt **648** indicates that the CRAM command **646** was successful and that Alice is now authenticated as Alice on Alice's server. While authenticated as Alice, a scan command **650** can result in output **652** that lists @-addresses on Alice's server, including @alice:myEmailPassword@alice,” “@barbara:location@alice,” “@barbara:phone@alice,” “@bob:phone@alice,” “public: email@alice,” “public:location@alice,” and “public:  custom charactercity@alice”.

(227) FIG. 7 illustrates a session **700** that includes use of an example profile and an example addressing scheme. A message **702** indicates that a connection has been made to a server for a user @alice. After an authentication process (not shown), a prompt **704** indicates that a user @bob has authenticated to the server for @alice. In response to submission of a “lookup:profile@alice” command **706**, an output **708** of “data:atsign://friends.profile@alice” is displayed.

(228) The “atsign://friends.profile@alice” value in the output **708** is an example of a URL (Uniform Resource Locator) pattern that includes an “atsign://” scheme **710** prepended to a URI of “friends.profile@alice”. The scheme **710** can identify an @ protocol for further parsing and resolution of the friends.profile@alice URI. The URI can be provided to a protocol-enabled application, for example, and the application can know how to parse the URI to further resolve friends.profile@alice. Schemes and URIs are discussed in more detail below.

(229) In the session **700**, a lookup command **712** is entered to resolve the friends.profile@alice URI. An output **714** is provided that includes contents of a friends.profile profile that the @alice user maintains (e.g. for close acquaintances). A profile can represent a persona, and can be a collection of information, shared with certain designated users. The friends.profile profile for @alice can include various items of information that the @alice user desires to share with a group of people, such as a first name, a phone number, social networking user names or addresses, and one or more email addresses. A user can configure different personas, by defining different profiles and sharing (e.g., enabling access to) the different profiles to different users. The @alice user may define a different profile for family or for work acquaintances, for example. While personas are described, profiles, or collections of information organized under a given address, can be used for other organizational purposes. A collection of information can include multiple items of information, separated by commas, as shown.

(230) The friends.profile profile includes a URL of “atsign://workphone@alice”, which is another

example of a URL that includes an atsign scheme. The `workphone@alice` URI can be further resolved, by looking up the `workphone@alice` URI on the server of `@alice`, to “mailto:alice@work.com”, for example. In general, use of the atsign scheme can enable reference pointers that point to other information and can result in the chaining together of responses when a given `service@atsign` pattern is resolved. For instance, an “atsign://email@alice” URI might resolve, for a given request, to “atsign://email.work@alice,” which may further (and eventually) resolve to “mailto:alice@work.com”. Chaining of responses can include use of URLs with different schemes in the responses so they can be parsed appropriately (e.g. “mailto://”, “https://”, or “atsign://” links, to name a few examples).

(231) An `atsign://URI` can resolve to a URL that points, for example, to image or video content that is stored in a server other than a user's secondary server. The URL or another value stored in user's personal directory can include a key that can be used to decrypt encrypted externally-stored content, for example.

(232) FIG. 8 is a block diagram illustrating an example system **800** for a private lookup transaction. The system **800** is similar to the system **300**, but the system **800** includes a different approach performed by a user's protocol server in contacting another protocol server on behalf of the user (or the user's application). In the example of FIG. 8, a user Bob **802** with a handle of `@bob` uses a device **804** to initiate an authenticated lookup of a `location@alice` address.

(233) In a first stage (e.g., denoted with a “1.”), a resolver **806** on the device **804** sends a query of “bob” to a root server **808**, to obtain server information for a protocol server for the `@bob` handle. In a second stage, the root server **808** returns server information having a pattern of “server2:port”—e.g., a server and port number of a protocol server **810** that is mapped to the `@bob` handle.

(234) In a third stage, the resolver **806** sends a “from:@bob” query to the protocol server **810**. In a fourth stage, the protocol server **810** responds to the “from:@bob” query with an authentication challenge. In a fifth stage, the resolver **806** (or another program, or Bob using another program) generates a signed challenge response. The signed challenge response is sent by the resolver **806** to the protocol server **810**. In a sixth stage, the protocol server **810** tests the signed challenge response, to determine whether the signed challenge response is valid with respect to the challenge sent to the resolver **806** in the fourth stage.

(235) In a seventh stage, in response to determining that the signed challenge response is valid, the protocol server **810** sends an indication (e.g., “@bob@”) to the resolver **806** to indicate that the resolver **806** has been authenticated to the protocol server **810** (e.g., as `@bob`). The “@bob@” indicator can correspond to a changing of a connection prompt for a connection between the resolver **806** and the protocol server **810**. A connection prompt of “@” can indicate an unauthenticated connection and a connection prompt of “@bob@” can indicate a connection authenticated as the user `@bob`.

(236) In an eighth stage, the resolver **806** sends a “lookup:location@alice” query to the protocol server **810**, using the authenticated connection. The protocol server **810** can perform a look up of the “location@alice” address on behalf of the resolver **806** (e.g., on behalf of `@bob`). The protocol server **810** may first need to obtain server information for a protocol server **812** that is mapped to the `@alice` handle.

(237) In a ninth stage, the protocol server **810** sends a query of “alice” to the root server **808** (or to another root server). In a tenth stage, the root server **808** (or another root server to which the “alice” query was sent) sends server information for the protocol server **812** (e.g., of a pattern of “server1:port”) to the protocol server **810**, in response to the “alice” query.

(238) In an eleventh stage, the protocol server **810** sends a “from:@bob” message to the protocol server **812**. In a twelfth stage, the protocol server **812** responds to the “from:@bob” request by sending a challenge to the protocol server **810**. The challenge is of a form of “xxx@bob:yyy”, which represents a challenge/request for the protocol server **810** to set an “xxx@bob” address to a

value of “yyy”.

(239) In a thirteenth stage, in response to the challenge from the protocol server **812**, the protocol server **810** updates, on the protocol server **810**, an “xxx@bob” address to the value of “yyy”. In a fourteenth stage, the protocol server **810** sends a proof-of-life (e.g., “pol”) message to the protocol server **812**, to indicate to the protocol server **812** that the requested challenge has been performed. The protocol server **810** is claiming that the protocol server **810** has successfully performed the requested challenge, as the @bob user.

(240) In a fifteenth stage, the protocol server **812** begins a process of validating whether the protocol server **810** successfully performed the requested challenge. The protocol server **812** can perform a public lookup on the protocol server **810** to determine whether the requested address was updated with the requested value. The protocol server **812** may first need to obtain server information for the protocol server **810** before performing the lookup. For example, in the fifteenth stage, the protocol server **812** sends a request to look up server information for the @bob user to the root server **808**. In a sixteenth stage, the root server **808** returns server information for the protocol server **810** (e.g., of a pattern of “server2:port”) to the protocol server **812**.

(241) In a seventeenth stage, the protocol server **812** establishes a separate connection with the protocol server **810** (e.g., different from the connection used in the eleventh, twelfth, thirteenth, and fourteenth stages). The protocol server **812** sends a “lookup:xxx@bob” request, over the separate connection, to the protocol server **810**. In an eighteenth stage, the protocol server **810**, in response to the “lookup:xxx@bob” request, retrieves a “yyy” stored at the “xxx@bob” address and sends a “data:yyy” response to the protocol server **812**.

(242) In a nineteenth stage, the protocol server **812** performs a test to determine whether the “data:yyy” value received from the protocol server **810** matches a requested “yyy” value that was included in the challenge sent to the protocol server **810** in the twelfth stage. If the test fails, the protocol server **812** can drop connection(s) to the protocol server **810**. If the test succeeds, the protocol server **812** can, in a twentieth stage, return an indication (e.g., “@bob@”) to the protocol server **810** over the first connection to indicate to the protocol server **810** that the protocol server **810** has been authenticated as the @bob user on the protocol server **812**.

(243) In a twenty-first stage, the protocol server **810** sends a “lookup:location@alice” request to the protocol server **812** as a request that has been authenticated as corresponding to the @bob user. In a twenty-second stage, the protocol server **812** retrieves a location@alice value (e.g., “SF Bay Area”) that has been specifically shared with the @bob user. The protocol server **812** sends a “data:SF Bay Area” response to the protocol server **810**. In a twenty-third stage, the protocol server **810** forwards the “data:SF Bay Area” response to the resolver **806**, for example, for presentation to the user Bob **802** and/or for use by an application on the device **804**.

(244) FIGS. **9A** and **9B** are a combined swim lane diagram illustrating an example method **900** for a private lookup transaction. The method **900** is similar to the method **400**, but method **900** uses a different approach of a user's protocol server contacting another protocol server, on behalf of the user (or the user's application). In the example of FIGS. **9A-9B**, a user Bob with a handle of @bob can use a device **902** to initiate an authenticated lookup of a location@alice address.

(245) At **904**, the device **902** sends a query of “bob” **906** to a root server **908**, to obtain server information for a protocol server **909** for the @bob handle.

(246) At **910**, the root server **908** returns server and port information **912** of “bobserver:6464” to the device **902**, in response to the query **906**.

(247) At **914**, the device **902** sends a “from:@bob” query **916** to the protocol server **909**.

(248) At **918**, the protocol server **909** responds to the “from:@bob” query **916** with an authentication challenge **920**. The authentication challenge **920** is of a form of “key:xxx@bob:yyy”.

(249) At **922**, the device **902** generates a signed challenge response **924** (e.g., using a PKAM command) and sends the signed challenge response **924** to the protocol server **909**.

(250) At **926**, the protocol server **909** tests the signed challenge response, to determine whether the signed challenge response **924** is valid with respect to the challenge **918**.

(251) At **928**, in response to determining that the signed challenge response **924** is valid, the protocol server **909** sends an indication **930** (e.g., “@bob@”) to the device **902** to indicate that the device **902** has been authenticated to the protocol server **909** (e.g., as @bob).

(252) At **932**, the device **902** sends a “lookup:location@alice” query **934** to the protocol server **909**, using the authenticated connection. The protocol server **909** can perform a look up of the “location@alice” address on behalf of the device **902** (e.g., on behalf of @bob). The protocol server **909** may first need to obtain server information for a protocol server **935** that is mapped to the @alice handle.

(253) At **936**, the protocol server **909** sends a query **938** of “alice” to the root server **908**.

(254) At **940**, the root server **908** sends server information **942** (e.g., of “aliceserver.com:6455”) for the protocol server **935** to the protocol server **909**, in response to the “alice” query **938**.

(255) At **944**, the protocol server **909** sends a “from:@bob” message **946** to the protocol server **935**.

(256) At **948**, the protocol server **935** responds to the “from:@bob” message **946** by sending a challenge **950** to the protocol server **909**. The challenge **950** has a pattern of “xxx@bob:yyy”, which represents a challenge/request for the protocol server **909** to set an “xxx@bob” address to a value of “yyy”.

(257) Referring now to FIG. **9B**, at **952**, in response to the challenge **950** from the protocol server **935**, the protocol server **909** updates, on the protocol server **909**, an “xxx@bob” address to the value of “yyy”.

(258) At **954**, the protocol server **909** sends a proof-of-life (e.g., “pol”) message **956** to the protocol server **935**, to indicate to the protocol server **935** that the requested challenge has been performed. The protocol server **909** is claiming that the protocol server **909** has successfully performed the requested challenge, as the @bob user.

(259) The protocol server **935** can begin a process of validating whether the protocol server **909** performed the requested challenge. The protocol server **935** can perform a public lookup on the protocol server **909** to determine whether the requested address was updated with the requested value. The protocol server **935** may first need to obtain server information for the protocol server **909**. For example, at **958**, the protocol server **935** sends a request **960** to look up server information for the @bob user to the root server **908**.

(260) At **962**, the root server **908** returns server and port information **964** (e.g., “bobserver:6464”) for the protocol server **909** to the protocol server **935**.

(261) At **966**, in a new session **968**, the protocol server **935** sends a “lookup:xxx@bob” request **970** to the protocol server **909**.

(262) At **972**, the protocol server **909**, in response to the “lookup:xxx@bob” request **970**, retrieves a “yyy” value stored at the “xxx@bob” address and sends a “data:yyy” response **974** to the protocol server **935**.

(263) At **976**, the protocol server **935** performs a test to determine whether the “data:yyy” value received from the protocol server **909** matches a requested “yyy” value that was included in the challenge **950**.

(264) At **978**, in response to the test succeeding, the protocol server **935** can return an indication **980** (e.g., “@bob@”) to the protocol server **909** to indicate to the protocol server **909** that the protocol server **909** has been authenticated as the @bob user on the protocol server **935**.

(265) At **982**, the protocol server **909** sends a “lookup:location@alice” request **984** to the protocol server **935** as a request that has been authenticated as corresponding to the @bob user.

(266) At **986**, the protocol server **935** retrieves a location@alice value (e.g., “SF Bay Area”) that has been specifically shared with the @bob user and sends a “data:SF Bay Area” response **988** to the protocol server **909**.

(267) At **990**, the protocol server **909** send a “data:SF Bay Area” value **992** to the device **902**, for example, for presentation to the user Bob and/or for use by an application on the device **902**.

(268) FIG. **10** is a block diagram illustrating an example system **1000** for configuring a device for protocol exchanges. In a first stage of a configuration process, a user Bob **1002** uses, for example, a web browser **1004** to connect to a registrar web server **1005** that provides a registrar website **1006**. The user Bob **1002** can use the registrar website **1006** to purchase an @sign handle (e.g., @bob). In a second stage, the user Bob **1002** uses the registrar website **1006** to select enabling of a device **1008** (e.g., @bob's device) for using the protocol.

(269) In a third stage, the registrar web server **1005** provides a code (e.g., a QR code) for display in the web browser **1004**. The code can include a CRAM shared secret mapped to the @bob handle. The registrar web server **1005** (or another server or system) can provide the shared secret to a protocol server **1010** that can handle protocol requests for the @bob handle. A protocol service can be initiated on the protocol server **1010** (e.g., at a particular port). The registrar web server **1005** (or another server or system) can provide server/port information for the protocol service running on the protocol server **1010** to a root server **1012**, for creation of a root record that maps the @bob handle to the protocol server **1010**.

(270) In a fourth stage, the user Bob **1002** uses the device **1008** to scan the presented QR code. The shared secret included in the QR code can be stored on the device **1008**, such as in a hardware trusted root.

(271) In a fifth stage, a settings application **1014** running on the device **1008** sends a “bob” query to the root server **1012**, to obtain protocol server information for the @bob handle. In a sixth stage, the root server **1012** sends server information (e.g., “server1:port”) for the protocol server **1010** to the settings application **1014**.

(272) In a seventh stage, the settings application **1014** sends a “from:@bob” message to the protocol server **1010**, to attempt to authenticate to the protocol server **1010**. In an eighth stage, the protocol server **1010** generates and sends key challenge information to the settings application **1014**. In a ninth stage, the settings application **1014** uses the shared secret and the key challenge to generate a digest. The digest can be included in a CRAM message that is sent to the protocol server **1010** in response to the key challenge.

(273) In a tenth stage, the protocol server **1010** determines whether the CRAM message (and digest) are correct, and in response to determining that the key challenge has been successfully met, provides an indicator (e.g., “@bob@”) to the settings application **1014** to communicate that the settings application **1014** has been successfully authenticated to the protocol server **1010** as @bob.

(274) In some implementations, the device **1008** uses a CRAM approach for initial setup/bootstrapping, but then moves to a PKAM approach for post setup authentication for protocol exchanges. For instance, in an eleventh stage, the settings application **1014** generates a public/private key pair (e.g., a RSA (Rivest—Shamir—Adleman) public/private key pair). In a twelfth stage, the settings application **1014** sends an update command to store a public key of the private key/public key pair on the protocol server **1010**.

(275) In a thirteenth stage, the settings application **1014** sends another “from:@bob” message to the protocol server **1010**, in a new, second (separate) session (e.g., second connection) with the protocol server **1010**. In a fourteenth stage, the protocol server **1010** generates and sends, over the second connection, key challenge information to the settings application **1014**. In this case, rather than responding to the key challenge using a CRAM message, the settings application **1014** can use the private key of the public/private key pair using a PKAM approach. For example, in a fifteenth stage, the settings application **1014** uses a PKAM command to generate a signed challenge that is signed with the private key. The signed challenge is sent by the settings application **1014**, over the second connection, to the protocol server **1010**. In some implementations, the settings application **1014** adds a client identifier (e.g., a client UUID) to the challenge information and then signs the combination of the challenge and the client identifier. Adding the client identifier can prevent

malicious attacks (e.g., spoofing attacks, replay attacks).

(276) In a sixteenth stage, the protocol server **1010** tests the signed challenge by using the public key received from the settings application **1014**. The protocol server **1010** can extract the key challenge information from the combination of the challenge and the client identifier, for example. If the test fails, the protocol server **1010** can drop the connection to the settings application **1014**. If the test succeeds, the protocol server **1010** can, in a seventeenth stage, return an indication (e.g., “@bob@”) to the settings application **1014**, over the second connection, to indicate to the settings application **1014** that the settings application **1014** has been authenticated as the @bob user on the protocol server **1010**. By receiving the successful authentication indication, the settings application **1014** can determine that the device **1008** has been configured properly for future protocol exchanges with the protocol server **1010**. In an eighteenth stage, the settings application **1014** can send an exit command, over the second connection, to the protocol server **1010**, as a request to close the second connection.

(277) In a nineteenth stage, the settings application **1014** can, over a first (e.g., still open) connection with the protocol server **1010**, send a delete command to delete the shared secret that was previously stored on the protocol server **1010**. The protocol server **1010** can delete the shared key from the protocol server **1010**, in response to the delete request. As mentioned, future authentication requests can use a PKAM approach and can be verified using the public key for @bob now stored on the protocol server **1010**.

(278) FIGS. **11A** and **11B** are a combined swim lane diagram illustrating an example method **1100** for configuring a device for protocol exchanges. At **1102**, a user Bob uses a device **1104** (e.g., a web browser running on the device **1104**) to connect to a registrar web server **1106** that provides a registrar website, to purchase an @sign handle (e.g., @bob).

(279) At **1108**, Bob uses the registrar website to select enabling of the device **1104** for using the protocol.

(280) At **1110**, the registrar web server **1106** generates a shared secret key that can be shared (e.g., for a bootstrap process) between the device **1104** and a protocol server **1112**.

(281) At **1114**, the registrar web server **1106** provides a code **1116** (e.g., a QR code) for display on the device **1104**. The code **1116** can include a shared secret mapped to the @bob handle.

(282) At **1118**, the registrar web server **1106** sends a request to the protocol server **1112** for the protocol server **1112** to store a secret key **1119** on the protocol server **1112**. A protocol service can be initiated on the protocol server **1112** (e.g., at a particular port).

(283) At **1120**, the registrar web server **1106** provides server/port information **1121** for the protocol server **1112** to a root server **1122**, for creation of a root record that maps the @bob handle to the protocol server **1112**.

(284) At **1124**, Bob uses the device **1104** to scan the presented QR code **1116** that was received from the registrar web server **1106**. The secret key included in the QR code **1116** can be stored on the device **1104**, such as in a hardware trusted root.

(285) At **1126**, the device **1104** sends a “bob” query **1128** to the root server **1122**, to obtain protocol server information for the @bob handle.

(286) At **1130**, the root server **1122** sends server information **1132** (e.g., in a “server:port” format) for the protocol server **1112** to the device **1104**.

(287) At **1134**, the device **1104** sends a “from:@bob” message **1136** to the protocol server **1112** in a first TCP/TLS session **1138** to attempt to authenticate to the protocol server **1112**.

(288) At **1140**, the protocol server **1112** generates and sends key challenge information **1142** to the device **1104**.

(289) At **1144**, the device **1104** uses the shared secret and the key challenge information **1142** to generate a CRAM message **1146** with an included digest that is sent to the protocol server **1112** in response to the key challenge information **1142**.

(290) At **1148**, the protocol server **1112** determines whether the CRAM message **1146** (and digest)

are correct, and, in response to determining that the key challenge has been successfully met, provides an indicator **1150** (e.g., “@bob@”) to the device **1104** to communicate that the device **1104** has been successfully authenticated to the protocol server **1112** as @bob.

(291) In some implementations, the device **1104** uses a CRAM approach for initial setup/bootstrapping but then moves to a PKAM approach for post setup authentication for protocol exchanges. Referring now to FIG. **11B**, at **1152**, the device **1104** generates a public/private key pair (e.g., a RSA public/private key pair), to be used for future PKAM operations.

(292) At **1154**, the device **1104** sends an update command **1155** to the protocol server **1112** to request storage of a public key of the private key/public key pair on the protocol server **1112**.

(293) At **1156**, the protocol server **1112** sends a success indicator **1157** to the device **1104** to indicate that the update command **1155** was processed successfully.

(294) At **1158**, the device **1104** sends, in a second TCP/TLS session **1160**, another “from:@bob” message **1162** to the protocol server **1112**.

(295) At **1164**, the protocol server **1112** generates and sends, in the second session, key challenge information **1166** to the device **1104**. In this case, rather than responding to the key challenge using a CRAM message, the device **1104** can use the private key of the public/private key pair using a PKAM approach.

(296) For example, at **1168**, the device **1104** uses a PKAM command to generate a signed challenge **1170** that is signed with the private key. The signed challenge **1170** is sent by the device **1104** to the protocol server **1112**.

(297) At **1172**, the protocol server **1112** tests the signed challenge **1170** by using the public key received from the device **1104**. If the test fails, the protocol server **1112** can drop the connection with the device **1104**. If the test succeeds, the protocol server **1112** can return an indication **1174** (e.g., “@bob@”) to the device **1104** to indicate that the device **1104** has been authenticated as the @bob user on the protocol server **1112**. By receiving the successful authentication indication, the device **1104** can determine that the device **1104** has been configured properly for future protocol exchanges with the protocol server **1112**.

(298) At **1176**, the device **1104** sends an exit command **1178** to the protocol server **1112**, as a request to end the second TCP/TLS session **1160**.

(299) At **1180**, the device **1104** sends a delete command **1181**, using a connection for the still open TCP/TLS session **1138**, to delete the shared secret that was previously stored on the protocol server **1112**. The protocol server **1112** can delete the shared key from the protocol server **1112**, in response to the delete command **1181**. The shared key can be removed from the protocol server **1112** so that a malicious entity that somehow accesses the protocol server **1112** can't obtain the shared secret from the protocol server **1112**. At **1182**, the protocol server **1112** sends an indication **1183** of a successful deletion to the device **1104**.

(300) At **1184**, the device **1104** sends an exit command **1186** to the protocol server **1112**, as a request to end the first TCP/TLS session **1138**.

(301) FIG. **12** is a flowchart of an example method **1200** for receiving and responding to a protocol request for user information. It will be understood that method **1200** and related methods may be performed, for example, by any suitable system, environment, software, and hardware, or a combination of systems, environments, software, and hardware, as appropriate. In some implementations, the method **1200** is executed by one or more components of the system **100** described above with respect to FIG. **1**. For example, the method **1200** can be executed by the settings application **145**, the protocol server **109**, or the user microservice **146** of FIG. **1**. As another example, the method **1200** can be executed by the protocol server **311** of FIG. **3** or the protocol server **413** of FIG. **4**.

(302) At **1202**, an authentication request is received, via a first connection and from a first requesting entity, to authenticate the first requesting entity to a first decentralized resource directory of a providing entity. The authentication request identifies the first requesting entity.

(303) At **1204**, an authentication challenge value and an authentication challenge key are generated, in response to the authentication request.

(304) At **1206**, an authentication challenge is provided, via the first connection and in response to the authentication request, to the first requesting entity for the first requesting entity to store the authentication challenge value for the authentication challenge key in a second decentralized resource directory of the first requesting entity.

(305) At **1208**, a confirmation is received, via the first connection and from the first requesting entity, that the authentication challenge value has been stored for the authentication challenge key in the second decentralized resource directory of the first requesting entity.

(306) At **1210**, a second connection is established with the second decentralized resource directory of the first requesting entity.

(307) At **1212**, a first lookup request is sent, via the second connection, for a value stored for the authentication challenge key in the second decentralized resource directory.

(308) At **1214**, a first response value is received, via the second connection and from the second decentralized resource directory, in response to the first lookup request.

(309) At **1216**, the first response value is compared to the authentication challenge value to determine whether the first response value matches the authentication challenge value.

(310) At **1218**, in response to determining that the first response value matches the authentication challenge value, the authentication request is responded to, via the first connection, with an indication that the first requesting entity is authenticated to the first decentralized resource directory. As another example, in response to determining that the first response value does not match the authentication challenge value, the first connection can be dropped.

(311) While not illustrated in FIG. **12**, it should be noted that a second lookup request can be received, via the first connection and from the first requesting entity, while the first requesting entity is authenticated to the first decentralized resource directory, for a value for a key stored in the first decentralized resource directory. A determination can be made that the first requesting entity has been previously provided access to a first stored value for the key. The first stored value for the key can be retrieved and provided as a second response value to the first requesting entity, in response to the second lookup request.

(312) Similarly, a third lookup request can be received, via an unauthenticated third connection and from an unauthenticated second requesting entity, for a value for the key stored in the first decentralized resource directory. A determination can be made that the key is publicly accessible. A publicly accessible value for the key can be retrieved and provided as a third response value, to the second requesting entity, in response to the third lookup request. The first stored value for the key provided to the authenticated first requesting entity can be different from the publicly accessible value for the key.

(313) In another example, a fourth lookup request can be received, via a third connection and from a third requesting entity while the third requesting entity is authenticated to the first decentralized resource directory, for a value for the key stored in the first decentralized resource directory. A determination can be made that the third requesting entity has been provided access to a third stored value for the key, where the third stored value is different from the first stored value and the publicly accessible value for the key. The third stored value can be retrieved and provided as a third response value to the third requesting entity, in response to the fourth lookup request.

(314) FIG. **13** is a flowchart of an example method **1300** for providing a protocol request for user information. It will be understood that method **1300** and related methods may be performed, for example, by any suitable system, environment, software, and hardware, or a combination of systems, environments, software, and hardware, as appropriate. In some implementations, the method **1300** is executed by one or more components of the system **100** described above with respect to FIG. **1**. For example, the method **1300** can be executed by the settings application **145**, the protocol server **109**, or the user microservice **146** of FIG. **1**. As another example, the method

1300 can be executed by the protocol server **314** of FIG. **3** or the protocol server **417** of FIG. **4**.

(315) At **1302**, an authentication request is sent, via a first connection, to authenticate as a requesting entity to a first decentralized resource directory of a providing entity. The authentication request identifies the requesting entity.

(316) At **1304**, an authentication challenge is received, via the first connection, from the providing entity, and in response to the authentication request, to store an authentication challenge value for a authentication challenge key in a second decentralized resource directory of the requesting entity.

(317) At **1306**, the authentication challenge value for the authentication challenge key is stored in the second decentralized resource directory.

(318) At **1308**, an authentication challenge response is sent to the providing entity requesting the providing entity to verify the authentication challenge.

(319) At **1310**, an indication is received from the providing entity that the requesting entity is authenticated to the first decentralized resource directory as the requesting entity.

(320) In some instances, a first lookup can be sent, while authenticated to the first decentralized resource directory and via the first connection, to the first decentralized resource directory for a value for a key stored in the first decentralized resource directory. A first value for the key can be received, in response to the first lookup request.

(321) As another example, a second lookup request can be sent, where the second lookup request is sent as an unauthenticated requesting entity and via a second connection to the first decentralized resource directory for the value for the key stored in the first decentralized resource directory. A second value for the key can be received, in response to the second lookup request, where the second value is different from the first value.

Onboarding and Initialization Options

(322) Various approaches can be used to onboard an application and/or a device for use with the @protocol. For example, QR codes, both onscreen and printed, can be used. As another example, private keys can be assigned per user or per application, and can be stored, for example, on the user's device or on a server. Further details are described in examples below.

(323) A user/device can be onboarded such that no entities other than the @sign user (e.g., entities such as the governing body **102**, the protocol provider server **103**, application developers **104**, root servers **105**, backup servers **106**, registrars **110**, or other entities) can access user data without the owner's permission. The system can be designed such that it is provably true that only the owner has control of their data. For example, encryption keys used by the decentralized resource directories are never stored in a place accessible to anyone other than the owner. For that purpose, the @protocol can use a hardware trusted root, typically found on the owner's personal devices (e.g. mobile phone) only accessible by the user. Such protections can be also combined with device-level security. The @protocol can also support backing up and synchronizing keys across multiple devices of a user.

(324) The initialization/bootstrapping of keys for decentralized resource directories and, more specifically, for secure key/value persistence can include a mechanism for generating an encryption key upon startup so that the encryption key is not stored in a central location. An encryption key can be kept in an owner's directory, but can be itself encrypted with the public key of a requester by the owner of the underlying data in response to the owner determining to accept a data request. In summary, the system can be designed such that any data stored in an owner's directory is impossible to see in an unencrypted form unless explicit permission has been granted by the data owner. Initialization and use of encryption keys is described in more detail below.

(325) FIG. **14** is a flow chart of an example onboarding process **1400**. In step **1.1**, a namespace API **1402** is used (e.g., by a registrar application) to generate a QR code **1404** and a shared secret for an @sign. In step **1.2**, the shared secret is shared with a secondary server **1406**. In step **1.3**, the secondary server **1406** generates a hive encryption key. The hive encryption key can be used when storing data in a file system **1408** on the secondary server **1406**. In step **1.4**, the secondary server

1406 stores the shared secret in a keystore **1410**.

(326) In step **2**, a user uses a first application **1411** on a user device **1412** to scan a presentation of the QR code **1404**. In step **3.1**, the first application **1411** generates a digest using the shared secret and sends the digest to the secondary server **1406**. The digest can be made using a CRAM command, for example. In step **3.2**, the secondary server **1406** sends a CRAM response message to the device **1412**.

(327) In step **4**, the first application **1411** can determine whether the CRAM response indicates authentication success. In step **5**, the first application **1411** generates a private/public key pair. In step **5.1**, the device **1412** sends the public key to the secondary server **1406**. In step **5.2**, the secondary server stores the public key in a keystore **1414**. The keystore **1414** can be a same or different keystore than the keystore **1410**. In step **5.3**, the first application **1411** generates a digest using the private key (e.g., using a pkam command) and sends the digest to the secondary server **1406**. In step **5.4**, the secondary server **1406** sends a PKAM response to the device **1412**.

(328) In step **6.1**, the first application **1411** generates a QR code **1415** that includes the public/private key pair. In step **6.2**, the QR code is sent to a printer **1416** for printing. In step **7.1**, the device **1412** sends a request to the secondary server **1406** for deletion of the shared secret. In step **7.2**, the secondary server deletes the shared secret from the keystore **1410**, in response to the deletion request. As indicated by a note **1418**, at the end of the process **1400**, the first application **1411** is onboarded on the device **1412**. A printout of the QR code **1415** can be used to onboard another application to the device **1412**, as described below with respect to FIG. **15**.

(329) FIG. **15** is a flow chart of an example onboarding process **1500** for onboarding a second application **1502**, e.g., on a device **1504**, for example after onboarding a first application (e.g., the first application **1411** described above with respect to FIG. **14**). In step **1**, the user uses the second application **1502** to scan a printed QR code **1506** (e.g., a code generated during an onboarding process for a first application). In step **2**, the second application **1502** generates a digest using a private key of a private/public key pair stored on the device **1504**, e.g., using a PKAM command. The digest is sent to a secondary server **1508**. In step **3**, the secondary server **1508** sends a PKAM response to the device **1504**. The second application **1502** can determine whether the PKAM response indicates authentication success. If the PKAM response indicates authentication success, the second application **1502** can be successfully onboarded on the device **1504**, as indicated by a note **1510**.

(330) A notes area **1512** lists advantages of using a paper QR code for onboarding additional applications to a device. For example, advantages can include 1) that a private key is accessible only by the user who has the paper QR code, 2) onboarding second (and additional) applications is a simple, shorter process (as compared to a first onboarding process), 3) subsequent applications can be onboarded to other devices of the user, and 4) if a device is reset or lost, a paper QR code can still be used to onboard an application (e.g., on a next device). A notes area **1514** lists potential considerations for using a paper QR code approach, as compared to alternative approaches described herein. For instance, the user may lose the paper version. As another example, while onboarding subsequent applications can involve a simple QR code scan, a first onboarded application may be required to perform additional steps, such as scanning a shared secret QR code and performing, for example, a CRAM operation.

(331) FIG. **16** is a flow chart of an example onboarding process **1600** using a separate private key per application. In step **1.1**, a namespace API **1602** is used (e.g., by a registrar application) to generate a QR code **1604** and a shared secret for an @sign. In step **1.2**, the shared secret is shared with a secondary server **1606**. In step **1.3**, the secondary server **1606** generates a hive encryption key for use when storing data in a file system **1608** on the secondary server **1606**. In step **1.4**, the secondary server **1606** stores the shared secret in a keystore **1610**.

(332) In step **2**, a user uses a first application **1611** on a user device **1612** to scan a presentation of the QR code **1604**. In step **3.1**, the first application **1611** generates a digest using the shared secret

and sends the digest to the secondary server **1606**. The digest can be made using a CRAM command, for example. In step **3.2**, the secondary server **1606** sends a CRAM response message to the device **1612**.

(333) In step **4**, the first application **1611** can determine whether the CRAM response indicates authentication success. In step **5**, the first application **1611** generates a first private/public key pair that includes a first public key and a first private key for the first application **1611**. In step **5.1**, the device **1612** sends the first public key to the secondary server **1606**. In step **5.2**, the secondary server **1606** stores the first public key in a keystore **1614**. The keystore **1614** can be a same or different keystore than the keystore **1610**. In step **5.3**, the first application **1611** generates a first digest using the first private key (e.g., using a pkam command) and sends the first digest to the secondary server **1606**. In step **5.4**, the secondary server **1606** sends a PKAM response to the device **1612**.

(334) In step **6**, the device **1612** determines whether the PKAM response indicates authentication success. In step **7.1**, the device **1612** sends a request to the secondary server **1606** for deletion of the shared secret. In step **7.2**, the secondary server **1606** deletes the shared secret from the keystore **1610**, in response to the deletion request. As indicated by a note **1616**, at the end of the process **1600**, the first application **1611** is onboarded on the device **1612**. Subsequent applications can be onboarded using a specific public/private key pair for each application.

(335) FIG. **17** is a flow chart of an example onboarding process **1700** for a second application using a separate private key per application. The process **1700** can be performed for a second application **1702** on a device **1704**. In step **1**, the second application **1702** generates a second private/public key pair that includes a second public key and a second private key for the second application **1702**. In step **1.1**, the device **1704** sends the second public key to a secondary server **1706**. In step **1.2**, the secondary server **1706** stores the second public key in a keystore **1708**. The keystore **1708** can be a same or different keystore than a keystore used to store a first public key (e.g., for a first application). In step **1.3**, the second application **1702** generates a second digest (e.g., different from a first digest created by a first application) using the second private key (e.g., using a pkam command) and sends the second digest to the secondary server **1706**. In step **1.4**, the secondary server **1706** sends a PKAM response to the device **1704**.

(336) In step **2**, the device **1702** determines whether the PKAM response indicates authentication success. As indicated by a note **1710**, at the end of the process **1700**, the second application **1702** is onboarded on the device **1704**. Subsequent applications, after the first and second applications, can be onboarded using a specific public/private key pair for each application. Advantages of using a private key per application, as compared to having a second application scan a QR code printed by a first application, can include avoiding a need to maintain or use a printed QR code. Additionally, subsequent applications for a user after a first application can be onboarded onto any device of the user. Considerations for using the private key per application approach can include considering an overhead of maintaining a private key per application. Also, if a user loses a device, the user may no longer have access to the private key stored on the device. The first onboarded application, with the private key per application approach, can have a capability to scan a shared secret and perform a CRAM method.

(337) FIG. **18** is a flow chart of an example onboarding process **1800** using a server-based private key. In step **1.1**, a namespace API **1802** is used (e.g., by a registrar application) to generate a QR code **1804** and a shared secret for an @sign. In step **1.2**, the shared secret is shared with a secondary server **1806**. In step **1.3**, the secondary server **1806** generates a hive encryption key for use when storing data in a file system **1808** on the secondary server **1806**. In step **1.4**, the secondary server **1806** stores the shared secret in a keystore **1810**.

(338) In step **2**, a user uses a first application **1811** on a user device **1812** to scan a presentation of the QR code **1804**. In step **3.1**, the first application **1811** generates a digest using the shared secret and sends the digest to the secondary server **1806**. The digest can be made using a CRAM

command, for example. In step 3.2, the secondary server **1806** sends a CRAM response message to the device **1812**. The first application **1811** can determine whether the CRAM response indicates authentication success. In step 3.3, the secondary server **1806** generates a first private/public key pair that includes a first public key and a first private key for the first application, and saves the first public key and the first private key in a keystore **1814**. The keystore **1814** can be a same or different keystore than the keystore **1810**.

(339) In step 4, the device **1812** sends a pkam request to retrieve the first private key to the secondary server **1806**. In step 4.1, the secondary server **1806** verifies the pkam request by determining whether the sender of the pkam request and an owner of the stored first private key correspond to the same @sign. In step 4.3, in response to a successful verification, the secondary server **1806** sends the first private key to the device **1812**.

(340) In step 5.1, the first application **1811** generates a first digest, e.g., using a pkam command, using the received first private key, and sends the first digest to the secondary server **1806**. In step 5.2, the secondary server **1806** sends a pkam response to the device **1812**. The first application can determine whether the pkam response indicates an authentication success. In step 6.1, the device **1812** sends a request to the secondary server **1806** for the secondary server **1806** to delete the shared secret. In step 6.2, the secondary server **1806** deletes the shared secret from the keystore **1810**, in response to the deletion request. As indicated by a note **1816**, at the end of the process **1800**, the first application **1811** is onboarded on the device **1812**.

(341) FIG. 19 is a flow chart of an example onboarding process **1900** for a second application using a server-based private key. In step 1, a second application **1902** on a device **1904** (e.g., different from a first application already onboarded onto the device **1904**) retrieves a second private key from a secondary server **1906** (e.g., the second private key being different from a first private key previously retrieved by the first application). For example, in step 1.1, the device **1904** sends a pkam request to retrieve the second private key to the secondary server **1906**. In step 1.2, the secondary server **1906** verifies the pkam request by determining whether the sender of the pkam request and an owner of the stored second private key correspond to the same @sign. In step 1.3, in response to a successful verification, the secondary server **1906** sends the second private key to the device **1904**.

(342) In step 2.1, the first application **1911** generates a second digest, e.g., using a pkam command, using the received second private key, and sends the second digest to the secondary server **1906**. The second digest is different from a first digest generated by the first application. In step 2.2, the secondary server **1906** sends a pkam response to the device **1912**. The second application **1902** can determine whether the pkam response indicates an authentication success. As indicated by a note **1910**, at the end of the process **1900**, the second application **1902** is onboarded on the device **1912**. Advantages of the server-based private key approach can be that a printed QR code does not need to be created or maintained and respective applications don't need to generate a private key.

(343) FIG. 20A illustrates an example system **2000** for onboarding a device/application. In step 1.1, a shared key is generated for an @sign and is embedded into a QR code **2002**. In step 1.2, a pkam key pair **2004** and an encryption key pair **2006** are generated. In step 2, the pkam key pair **2004** and the encryption key pair **2006** are encrypted using the shared key as encrypted key pairs, and stored in a key pairs storage **2008**. In step 3.1, the encrypted key pairs are streamed to a user device **2010**. In step 3.2, the QR code **2002** is scanned by the user device **2010**, thereby enabling the user device **2010** to obtain the shared key. In step 4, the user device **2010** decrypts the key pairs received in step 3.1 using the shared key obtained in step 3.2. The decrypted key pairs can be stored on the user device **2010**.

(344) FIG. 20B illustrates an example system **2020** for onboarding a device/application using a registrar. A user may wish to onboard a settings application **2022** on a device **2024** of the user. The settings application **2022** can communicate with a registrar **2025** using an API **2026** provided by the registrar **2025**. For example, in a first step, the settings application **2022** sends a request to the

registrar **2025** (e.g., using the API **2026**) for a list of available @signs. In a second step, the registrar **2025** sends a list of available @signs to the device **2024** for presentation in the settings application **2022**. As another example, the settings application **2022** can query the registrar **2025** regarding availability of a specific @sign, rather than request a list of available @signs. The user can select an available @sign, based on information or a response received from the registrar **2025**. The user can select an option in the settings application **2022** to onboard the settings application **2022** using an email or SMS (Short Message Service) verification, rather than other approaches such as QR code scanning.

(345) In a third step, the settings application **2022** sends a selected @sign and a user phone number or user email address to the registrar **2025**. As a first option for a fourth step, if the registrar **2025** receives a user phone number from the settings application **2022**, the registrar **2025** can generate an OTP (One Time Password) and provide the OTP to a SMS gateway **2028**. The SMS gateway **2028** can be used to send a SMS message with the OTP to the device **2024**. The SMS message with the OTP can be displayed on the device **2024** using a SMS application **2030**. As a second option for the fourth step, if the registrar **2025** receives a user email from the settings application **2022**, the registrar **2025** can generate an OTP and provide the OTP to an email gateway **2032**. The email gateway **2032** can be used to send an email message with the OTP to the device **2024**. The email message with the OTP can be displayed on the device **2024** using an email application **2034**.

(346) In a fifth step, the user views the OTP using either the SMS application **2030** or the email application **2034** and then enters the OTP into the settings application **2022**. In a sixth step, the settings application **2022** sends the OTP to the registrar **2025**. In a seventh step, the registrar **2025** sends a CRAM shared secret, for authentication to the secondary server of the user, to the settings application **2022**. The shared secret can be used as described above (e.g., to create and send a digest to the secondary server of the user).

(347) FIG. **20C** is a flow chart of an example onboarding process **2050** for onboarding a device/application using a registrar. A user may wish to onboard a settings application **2052** on a device of the user. The settings application **2052** can communicate with a registrar **2055** using an API **2056** provided by the registrar **2055**. For example, at **2058**, the settings application **2052** sends a request to the registrar **2055** (e.g., using the API **2026**) for a list of available @signs. At **2060**, the registrar **2055** sends a list of available @signs to the device **2024** for presentation in the settings application **2052**. As another example, the settings application **2052** can query the registrar **2055** regarding availability of a specific @sign, rather than request a list of available @signs. The user can select an available @sign, based on information or a response received from the registrar **2055**. The user can select an option in the settings application **2052** to onboard the settings application **2052** using an email or SMS (Short Message Service) verification, rather than other approaches such as QR code scanning.

(348) At **2062**, the settings application **2052** sends a selected @sign and a user phone number or user email address to the registrar **2055**. If the registrar **2055** receives a user phone number from the settings application **2052**, the registrar **2055** can, at **2064**, generate an OTP and provide the OTP to a SMS gateway **2066**. The SMS gateway **2066** can be used to send a SMS message with the OTP to the device of the user. The SMS message with the OTP can be displayed on the device using a SMS application **2068**. As another example, if the registrar **2055** receives a user email from the settings application **2052**, the registrar **2055** can, at **2070**, generate an OTP and provide the OTP to an email gateway **2072**. The email gateway **2072** can be used to send an email message with the OTP to the device of the user. The email message with the OTP can be displayed on the device using an email application **2074**.

(349) At **2076**, if the device receives the OTP from the email gateway **2072**, the user views the OTP using the email application **2074** and then enters the OTP into the settings application **2052**. As another example, at **2078**, if the device receives the OTP using the SMS gateway **2066**, the user views the OTP using the SMS application **2068** and then enters the OTP into the settings

application **2052**. At **2080**, the settings application **2052** sends the OTP to the registrar **2055**. At **2082**, the registrar **2055** sends a CRAM shared secret, for authentication to the secondary server of the user, to the settings application **2052**. The shared secret can be used as described above (e.g., to create and send a digest to the secondary server of the user).

(350) FIG. **21** is a flowchart of an example method **2100** for lifecycle management of a secondary server.

(351) At **2102**, the secondary server is started. The secondary server can be associated with a particular IP address (which may be mapped to a host name) and port, for example.

(352) At **2104**, a determination is made as to whether the host and port are valid. If either or both of the host and port are not valid, the method **2100** can end (e.g., and an administrator can be notified so as to troubleshoot host and/or port issues).

(353) At **2106**, if the host and port are valid, hive storage at the secondary server is initialized. For example, secret management tasks can be initiated, a storage area can be initialized, and various adapters for different types of data can be registered. As another example, one or more expiration tasks can be scheduled, such as for managing data that has a TTL setting. For instance, an expiration task can invalidate or delete expired data when TTL time points are reached.

(354) At **2108**, a determination is made as to whether all initialization steps have completed successfully. If not all initialization steps have been completed successfully, the method **2100** can end (e.g., and an administrator can be notified so as to troubleshoot initialization issues).

(355) At **2110**, if all initialization steps have been completed successfully, a server process can be started. For instance, a secure connection can be created and a listening task can be initiated.

(356) At **2112**, a post-start process can be initiated, which can include, for example, storing of a secret for the user associated with the secondary server.

(357) FIG. **22** is a flowchart of an example method **2200** for data expiration management.

(358) At **2202**, a query is performed for expired keys in a hive storage area of a secondary server. For example, keys that have a TTL timestamp less than a current timestamp can be identified.

(359) At **2204**, a determination is made as to whether any expired keys were returned from the query. If no expired keys were returned from the query, the method **2200** can end.

(360) At **2206**, if at least one expired key was returned from the query, then the returned expired key(s) can be deleted from the hive storage on the secondary server.

(361) At **2208**, one or more notifications can be sent to notify of deletion of the expired key(s). For example, one or more parties may have been subscribed to receive notices. As another example, a notification can be sent to a settings application (e.g., and/or one or more client applications) on a user device of a user associated with the secondary server.

(362) FIG. **23** illustrates an example system **2300** for saving a profile using encryption. In a first stage, a user **2302** installs a settings application **2304** on a user device **2306**. Installing the settings application **2304** can result in a local hive storage (e.g., local storage) **2308** and a SDK (Software Development Kit) being installed on the user device **2306**.

(363) In a second stage, the user **2302** uses the settings application **2304** to generate an encryption key and store the encryption key in an entry **2310** in the local hive storage **2308**.

(364) In a third stage, the user creates keys, for example, using one or more protocol applications. The keys can be saved to secondary storage **2311** on a secondary server **2312** (and, optionally, also to the local hive storage **2308**). For example, an entry **2314** in the secondary storage **2311** includes a key, a key value, an encryption key, and, optionally, a TTL value. The secondary server **2312** is associated with (e.g., known to) a root server **2316**.

(365) FIG. **24** illustrates a system for @sign activation. In a first stage (e.g., a circled “1”), in response to an @sign user providing an input on an @registrar web application **2402** (or other user interface) as a request to activate their @sign, the @registrar web application **2402** can start an @sign activation process for the @sign. The activation process can include an HTTPS call to a @namespace API **2404** that can be, for example, a POST command such as “POST:@sign, APIkey,

options”. Options may relate to location, provider, or other optional features.

(366) In a second stage, the @namespace API **2404** receives and validates the activation request received from the @registrar web application **2402**. For example, the @namespace API **2404** can validate formatting of the @sign and determine whether the @sign is reserved, banned, or already activated.

(367) As shown in FIG. **25**, banned @signs **2502** can include a list of offensive words that can be banned from use. Banned @signs **2502** are not available for purchase or use with the @protocol. Reserved @signs **2504** can be for trademarks that, while valid for use, may create infringement and can therefore be withheld except for use by an owner of the trademark. Registered @signs **2506** have been purchased (or otherwise reserved) but not activated, deployed or added to an @root server. Activated @signs **2508** have been purchased, activated, deployed and added to @root server(s).

(368) Returning to FIG. **24**, if the @sign is improperly formatted or is reserved, banned, or already activated, the @namespace API **2404** can return an error to the @registrar web application **2402**. The @namespace API **2404** can generate a CRAM shared secret, generate a QR code that includes the shared secret and the @sign, and generate a URL that can be used for retrieving a QR code image.

(369) In a third stage, the @namespace API **2404** calls an appropriate @infrastructure API to create a new @server. For example, the following POST command can be sent from the @namespace API to an infrastructure API component **2406**: “HTTPS POST:@sign, APIkey, CRAM shared secret”.

(370) In a fourth stage, the infrastructure API component **2406** receives and validates the activation request received from the @namespace API **2404**. In some implementations, the infrastructure API component **2406** can perform similar formatting or banned/reserved/active tests on the requested @sign as done by the @namespace API **2404**. Other validation operations can be performed.

(371) In a fifth stage, the infrastructure API component **2406** returns a validation response to the @namespace API **2404** that indicates validation success or failure. In response to validation success, the infrastructure API component **2406** can begin asynchronous deployment of an @server for the requested @sign (e.g., as described in more detail below for stages nine, ten, and eleven).

(372) In a sixth stage, the @namespace API **2404** returns a response to the @registrar web application **2402** regarding activation of the @sign. If validation steps have thus far been successful, the response sent to the @registrar web application **2402** can indicate validation success and can include QR image data that includes the shared secret and the requested @sign. If validation steps have not all been successful, the response sent to the @registrar web application **2402** can include an indication of validation failure.

(373) In a seventh stage, the @registrar web application **2402** can process the response received from the @namespace API **2404**, which can include displaying an error message in the case of an error response or displaying the QR code image (e.g., using a received image URL). The user can be prompted to scan the QR code using a settings application **2408** (e.g., an @me application).

(374) In an eighth stage, the settings application **2408** scans the QR code displayed in the @registrar web application **2402** and then begins a polling process of sending polling messages to an @root server **2410** to determine (based on a response from the @root server **2410**) when an @server **2412** has been successfully registered.

(375) As mentioned, stages nine, ten, and eleven may have been occurring, or may occur, asynchronously after successful validation by the infrastructure API component **2406**. For example, in a ninth stage, the infrastructure API component **2406** can create a new container for the @sign in the @sign server **2412** using the CRAM secret. In a tenth stage, the infrastructure API component **2406** can monitor the new @server **2412** to determine when the new @server **2412** is online. In an eleventh stage, the infrastructure API component **2406** can send an update message to the @root server **2410** indicating the online status of the @server **2412**.

(376) In a twelfth stage, the setting application **2408** can locate (e.g., by the polling process) an

@root entry for the @server **2412** at the root server **2410**. The settings application **2408** can also poll the @server **2412** to determine when the @server **2412** is ready for a CRAM to PKAM initialization.

(377) In a thirteenth stage, after the @server **2412** is ready for the CRAM to PKAM initialization, the settings application **2408** initializes a PKAM credentials process (e.g., creation of a private key/public key pair with the public key only on the @server **2412**). The CRAM shared key can be deleted by the @server **2412** after the PKAM credentials are configured.

(378) In a fifteenth stage, the @registrar web application **2402** can determine/track status of the activation by looking up and locating a root entry for the @server **2412** at the @root server **2410**.

(379) In a sixteenth stage, the @registrar web application **2402** can perform a scan of the @server **2412** to determine that the @server **2412** is active. Upon successful detection of an active @server **2412**, the @registrar web application **2402** can present an activation success message for the user.

(380) In some implementations, the @namespace API **2404** can periodically poll the @server **2412**, for example, to check to see if the @server **2412** is no longer online or if the CRAM shared secret is no longer stored at the @server **2412**. If the @server **2412** is no longer online or if the CRAM shared secret is no longer stored at the @server **2412**, QR code data can be deleted. In some implementations, the infrastructure API component **2406** can periodically check to see if the @server **2412** is still active or if an activation has failed to complete. In response to a failed activation or an inactive @server **2412**, the infrastructure API component **2406** can delete the @sign:location from the @root server **2410**, perform removal activities for the @server **2412**, and return a port that was used for the @server **2412** to an available pool of port numbers.

Streaming and File Transfers

(381) Some previously-developed applications can suffer from limitations of only working between devices of the same platform and/or only over a LAN connection. An @file_sharing application can be provided with/by the @protocol and can support file transfer without requiring e-mail or a mass storage device. The @protocol file sharing feature can be a secure, surveillance-free application that can work between all devices and locations by virtue of the following features/aspects of the @protocol: 1) an ability to locate a user's device using the @protocol namespace directory; 2) a permissioning system's ability to control access between the entities; 3) an ability to encrypt file contents; and 4) a use of P2P communications between devices that prevents surveillance.

(382) FIG. **26** illustrates a system **2600** that illustrates stages of protocol-enabled file sharing. In a first stage **2602**, a sender initiates a file transfer and a recipient acknowledges the transfer. In some implementations, some acknowledgements are automatically sent by a recipient. Transfer initiation can have a timeout, whereby a transfer is cancelled if a recipient does not respond. In some implementations, parallel transfers from multiple users to a same recipient are supported.

(383) In a second stage **2604**, a communication channel is established between the sender and the recipient. In some cases, a preconfigured connection can exist between a particular sender and recipient, and the preconfigured connection can be identified without needing to be re-established.

(384) In a third stage **2606**, the file is transferred from the sender to the recipient.

(385) In a fourth stage, housecleaning tasks are performed, such as marking a file as transferred. The file can be marked as transferred to prevent or warn about a duplicate transfer. In some cases, a file-complete notification and/or acknowledgement are communicated over the channel.

(386) In general, different approaches can be used for file transfers or general streaming content from one protocol user to another. For example, server hop(s), tokens, notifications, proof of life, and/or intermediate servers can be used. Streaming data, e.g., from an @bob user to an @alice user can pose various challenges. For example, each user may be at different locations (e.g., perhaps in different countries), at a same location, on a same network (e.g., same WiFi network), or on different networks. Streaming can involve using a device (either a user device or an intermediary device) that has a dedicated IP address. As another example, streaming can involve using user's

secondary servers. However, some secondary servers may be configured as micro servers (e.g., not specifically designed for streaming a large amount of data), so avoiding bottlenecks when transferring data at scale, using secondary servers, can be a point of consideration. Further details of streaming approaches are discussed below.

(387) FIG. 27 is a flowchart of example streaming method 2700. In preconfiguration steps 2702 and 2704 monitoring processes are established between an Alice client (e.g., for an @alice user) 2706 and Alice's secondary server 2708 and between Bob's client 2710 (e.g., for an @bob user) and Bob's secondary server 2712, respectively.

(388) In step 1, Alice's client 2706 sends a stream initialization message to Alice's server 2708 that includes a stream identifier. In step 2, Alice's server 2708 sends a notification message about the stream to Bob's secondary server 2712, that includes the stream identifier and a file length value. In step 3, Bob's server 2712 sends a notification message, which can be sent using a preconfigured monitoring process, to Bob's client 2710. In step 4, Bob's client 2710 sends an acknowledgement message to Alice's server 2708.

(389) In step 5, Alice's client 2706 receives the acknowledgement (which can be received using a preconfigured monitoring process). As noted in step 5.1, Alice's client 2706 can receive a timeout event if the acknowledgement is in fact not received. In step 6, after receiving the acknowledgement, Alice's client 2706 streams data for the file to Alice's secondary server 2708. In step 7, Alice's secondary server 2708 relays received streaming data to Bob's client 2710. In step 8, Bob's client 2710 saves the data (e.g., to a file system) when an EOF (end of file) is reached in the received streaming data. In step 9, Bob's client 2710 sends a stream complete message (that includes the stream identifier) to Alice's secondary server 2708. In step 10, Alice's secondary server 2708 sends marks the stream as completed and sends a stream complete message to Alice's client 2706. Alice's client 2706 can also mark the stream as completed.

(390) FIG. 28 is a flowchart of example streaming method 2800 using a server hop. In preconfiguration steps 2802 and 2804 monitoring processes are established between an Alice client (e.g., for an @alice user) 2806 and Alice's secondary server 2808 and between Bob's client 2810 (e.g., for an @bob user) and Bob's secondary server 2812, respectively.

(391) In step 1, Alice's client 2806 sends a stream initialization message to Alice's server 2808 that includes a stream identifier.

(392) In step 2, Alice's server 2808 sends a notification message about the stream to Bob's secondary server 2812, that includes the stream identifier and a file length value.

(393) In step 3, Bob's server 2812 sends a notification message, which can be sent using a preconfigured monitoring process, to Bob's client 2810.

(394) In step 4, Bob's client 2810 sends an acknowledgement message to Bob's secondary server 2812.

(395) In step 5, Bob's secondary server 2812 sends an acknowledgement message to Alice's secondary server 2808.

(396) In step 6, Alice's secondary server 2808 sends an acknowledgement message to Alice's client 2806. As noted in step 6.1, Alice's client 2806 can receive a timeout for the stream if an acknowledgement is in fact not received.

(397) In step 7, Alice's client 2806 streams data for the file to Alice's secondary server 2808.

(398) In step 8, Alice's secondary server 2808 relays stream data to Bob's secondary server 2812.

(399) In step 9, Bob's secondary server 2812 relays stream data to Bob's client 2810.

(400) In step 10, Bob's client 2810 sends a stream-done message to Bob's secondary server 2812.

(401) In step 11, Bob's secondary server 2812 sends a stream-done message to Alice's secondary server 2808.

(402) In step 12, Alice's secondary server 2808 sends a stream-done message to Alice's client 2806. Alice's client 2806 can mark the stream transfer as completed.

(403) FIG. 29 is a flowchart of example streaming method 2900 using a token. In preconfiguration

steps **2902** and **2904** monitoring processes are established between an Alice client (e.g., for an @alice user) **2906** and Alice's secondary server **2908** and between Bob's client **2910** (e.g., for an @bob user) and Bob's secondary server **2912**, respectively.

(404) In step **1**, Alice's client **2906** sends an update request to Alice's secondary server **2908** to update a stream token for an intended stream.

(405) In step **2**, Alice's secondary server **2908** sends a notification of the intended stream to Bob's secondary server **2912**.

(406) In step **3**, Bob's server **2912** sends a notification message, which can be sent using a preconfigured monitoring process, to Bob's client **2910**.

(407) In step **4**, Bob's client **2910** sends a lookup request to Bob's secondary server, for looking up the stream token for the intended stream. Bob's secondary server **2912** can look up the stream token, from Alice's secondary server **2908**.

(408) In step **5**, Bob's client **2910** sends a stream initialization message to Alice's secondary server **2908**.

(409) In step **6**, Alice's secondary server **2908** sends a notification to Alice's client **2906**.

(410) In step **7**, Alice's client streams data for the stream to Alice's secondary server **2908**.

(411) In step **8**, Alice's secondary server **2908** relays streamed data to Bob's client **2910**.

(412) In step **9**, Bob's client **2910** saves the streamed data, e.g., to a file system on Bob's client **2910**, when the stream transfer is complete.

(413) FIG. **30** is a flowchart of example streaming method **3000**. In preconfiguration steps **3002** and **3004** monitoring processes are established between an Alice client (e.g., for an @alice user) **3006** and Alice's secondary server **3008** and between Bob's client **3010** (e.g., for an @bob user) and Bob's secondary server **3012**, respectively.

(414) At step **1**, Alice's client **3006** sends a notification (including a particular notification identifier) of an intended stream, to Alice's secondary server **3008**.

(415) In step **2**, Alice's secondary server **3008** sends a notification to Bob's secondary server **3012**.

(416) In step **3**, Bob's secondary server **3012** sends a notification message, which can be sent using a preconfigured monitoring process, to Bob's client **3010**.

(417) In step **4**, Bob's client **3010** sends an acknowledgement (including the notification identifier) of the intended stream to Alice's secondary server **3008**.

(418) In step **5**, Alice's secondary server **3008** forwards the acknowledgement message to Alice's client (e.g., using the preconfigured monitoring process). As indicated in step **5.1**, Alice's client **3006** can receive a timeout for the intended stream if in fact an acknowledgment is not received.

(419) In step **6**, Alice's client **3006** streams data to Alice's secondary server **3008**.

(420) In step **7**, Alice's secondary server **3008** relays streamed data to Bob's client **3010**.

(421) In step **8**, Bob's client saves the streamed data, e.g., in a file system, when the end of the stream is reached.

(422) In step **9**, Bob's client **3010** sends a stream complete notification (including the notification identifier) to Alice's secondary server **3008**.

(423) In step **10**, Alice's secondary server **3008** sends the stream complete notification to Alice's client **3006**. The stream can be marked as complete on Alice's client **3006** and on Alice's secondary server, or tracking entries that include the notification identifier can be deleted (e.g., based on stream completion).

(424) FIG. **31** is a flowchart of example streaming method **3100** using a proof of life approach. In steps **3102** and **3104** monitoring processes are established between an Alice client (e.g., for an @alice user) **3106** and Alice's secondary server **3108** and between Bob's client **3110** (e.g., for an @bob user) and Bob's secondary server **3112**, respectively.

(425) In step **1**, Alice's client **3106** sends a stream initialization message to Alice's secondary server **3108**.

(426) In step **2**, Alice's secondary server **3108** forwards a stream initialization message to Bob's

secondary server **3112**.

(427) In step **3**, Bob's secondary server **3112** sends a from message to Alice's secondary server **3108** to initiate a pol challenge.

(428) In step **4**, Alice's secondary server **3108**, as part of the pol challenge, sends a proof challenge for Bob's secondary server to lookup a requested item on Bob's secondary server **3112** placed by Alice's secondary server **3108**.

(429) In step **5**, Bob's secondary server **3112** performs the requested lookup and sends the looked-up data, as proof for the pol challenge, to Alice's secondary server **3108**.

(430) In step **6**, in response to confirming the pol challenge, Alice's secondary server **3108** sends a notification of a server-to-server handshake success (e.g., the successful pol challenge) to Alice's client **3106**.

(431) In step **7**, Alice's client **3106** streams data to Alice's secondary server **3108**.

(432) In step **8**, Alice's secondary server **3108** relays streamed data to Bob's client **3110**. In step **9**, Bob's client saves the streamed data, e.g., in a file system, upon detecting an end of stream.

(433) FIG. **32** is a flowchart of example streaming method **3200** using an intermediate streaming server **3202**. The intermediate streaming server **3202** can be used, rather than a secondary server **3204** for an @alice user, when the @alice user desires to stream data to an @bob user.

(434) In step **1**, Alice's client **3206** sends a stream:to@bob streaming message to the intermediate streaming server **3202** regarding an intent to stream data to the @bob user.

(435) In step **2**, the intermediate streaming server **3202** sends a stream:from@alice message to Bob's secondary server **3208**.

(436) In step **3**, Bob's secondary server **3208** sends a notification message about the stream to Bob's client **3210**.

(437) In step **4**, Bob's client **3210** sends an acknowledgement to the intermediate streaming server **3202**.

(438) In step **5**, the intermediate streaming server **3202** forwards the acknowledgement to Alice's client **3206**.

(439) In step **6**, Alice's client **3206** sends streaming data to the intermediate streaming server **3202**.

(440) In step **7**, the intermediate streaming server **3202** relays streaming data to Bob's client **3210**.

(441) In step **8**, Bob's client sends a stream:done notification to the intermediate streaming server **3202** when Bob's client **3210** has finished receiving the streamed data.

(442) In step **9**, the intermediate streaming server **3202** forwards the stream:done notification to Alice's client **3206**.

Data Sharing and Encryption

(443) A major issue historically for the adoption of secure communication systems (e.g. PGP (Pretty Good Privacy)) has been the challenge of distributing signing keys used to encrypt/decrypt information exchanged between entities. The @protocol solves this problem in a unique manner by making it easy to securely share information by permission to some other entity. For example, "publickey@alice" can be a common syntax and approach for sharing a public key, and "atsign://publickey@alice" can be a fully qualified endpoint URL that can be distributed. While a public key used for PKI (Public Key Infrastructure) encryption can be shared publicly, typically a user may securely share a public key only to specified users. The @protocol supports both public and private sharing of encryption keys. The @protocol provides a simple means of creating and distributing PKI keys. Distribution and sharing of public keys can come as a side effect of the fact that the protocol itself relies on key sharing capability to securely share data with others. The result is an Internet-scale directory for billions of @sign owners to easily publish and lookup PKI keys on the Internet. Thus the @protocol provides a means for distributing and managing PKI keys at a massive scale.

(444) The @protocol can also provide a system for distributing and managing PKI keys that can provide different responses based on who is requesting a public key. For instance, in addition to

distributing encryption keys as outlined above, the protocol can determine which respective encryption key to provide among multiple encryption keys in response to a request, wherein the respective encryption key to provide depends on an identity of a requestor. That is, the protocol can support polymorphic sharing and distribution of encryption keys. This capability can provide additional assurance that only an intended recipient is able to access shared information.

(445) The @protocol's polymorphism capability can enable providing different responses to a query based on which @sign is requesting the data. Polymorphism can be applicable for public, as well as private, sharing. Polymorphism capabilities can have important implications for managing PKI keys and other types of data. For example, a protocol query such as `publickey@alice` can be responded to with a public key value that is shared without using the proof of life handshake, as a public value shared with all users or a response can include different values for a same `publickey@alice` query based on a "from" parameter as part of the pol handshake. A given user can share a different public key with each of multiple different other users, for example. For instance, @alice can share a first public key with @bob and a second public key with @sara, each in response to a lookup by @bob or @sara of a `publickey@alice` key, respectively, based on a determination by @alice for which public key to share with which user. Polymorphic responses to queries for keys can provide a simple, scalable system for managing public and private PKI keys.

(446) Additionally, the @protocol can provide a system for distributing content to multiple recipients that is encrypted one time such that only the intended recipients can access the content by using those intended recipients' respective personal keys. In further detail, the @server sharing protocol can support encryption of information that is encrypted once with a shared symmetric key, and is then shared with any number of others in such a manner that only the recipient can decrypt the information using their private key with asymmetric decryption. Such sending can be subject to other privacy controls inherent in the @protocol. For example, decrypted information can be set to remain accessible for some defined time period after which it will be automatically deleted.

(447) With the @protocol, a hybrid symmetric/asymmetric encryption scheme can be used wherein the information can be first encrypted using a shared symmetric key set by the sender. The shared symmetric key can then itself be encrypted using asymmetric encryption with the sender's private key and the recipient's public key from the recipient's @server. A result of the hybrid symmetric/asymmetric scheme can be that information that is shared by the sender can only be retrieved using the recipient's private key which can provide access to the shared symmetric key set by the sender. The recipient can then use the shared symmetric key to decrypt the information shared with the recipient.

(448) FIG. 33 illustrates an example system 3300 for use of encryption keys. As indicated by a note 3301, a net result of a security approach can be that only user-encrypted data is stored on secondary servers and only authorized @sign users can access data based on sharing privileges granted by data owners.

(449) As described in a note 3302, during setup for an @sign, a public/private key pair is created and a public key service is populated. For example, for an @alice user 3303, a key pair 3304 can be created on Alice's device 3304 and a `publickey@alice` value 3305 can be stored on Alice's secondary server 3306. As another example, for an @bob user 3307, a key pair 3308 can be created on Bob's device 3309 and a `publickey@bob` value 3310 can be stored on Bob's secondary server 3311.

(450) In a first stage 3312 of a use case scenario, Alice wants to share a `phone@alice` value 3313 with the @bob user 3307.

(451) In a second stage 3314, Alice's device 3304 generates a shared key (e.g., an Advanced Encryption Standard shared key) for the @bob user 3307.

(452) In a third stage 3316, Alice's device 3304 requests that Alice's secondary server 3306 perform a lookup for the public key 3310 for the @bob user 3307 that is stored on Bob's secondary server 3311. Alice's device 3304 encrypts the shared key with Bob's public key and requests storage of an

encrypted shared key **3318** on Alice's secondary server **3306**.

(453) In a fourth stage **3320**, Alice's device **3304** encrypts the shared key with Alice's public key **3305**, and stores an encrypted shared key **3322** on Alice's secondary server **3306**.

(454) In a fifth stage **3324** corresponding to Bob desiring to access Alice's phone number, Bob's device **3309** sends a lookup request to Bob's secondary server **3311** for Bob's secondary server **3311** to look up the shared encrypted key **3318**. Bob's device **3309** uses a private key of the key pair **3308** to decrypt the encrypted shared key.

(455) In a sixth stage **3326**, Bob's device **3309** sends a lookup request to Bob's secondary server **3311** for Bob's secondary server **3311** to look up the phone@alice value **3313** from Alice's secondary server **3306**. Bob's device **3309** receives the phone@alice value **3313** as an encrypted phone number. Bob's device **3309** uses the shared key obtained from decrypting the encrypted shared key to decrypt the encrypted phone number, to obtain an unencrypted phone number value (e.g., Alice's phone number she has shared with Bob).

(456) FIG. **34** illustrates an example system **3400** for use of encryption keys. In step **1.1**, a user Alice (e.g., corresponding to Alice's device **3402**) decides to share a phone@alice value with a user Bob (corresponding to Bob's device **3404**).

(457) For instance, in step **1.2**, Alice's device **3402** generates a shared key **3405** for the user Bob.

(458) In step **1.3**, Alice's device initiates a lookup of Bob's public key. Alice's device **3402** sends a lookup request **3406** to Alice's secondary server **3408**. Alice's secondary server **3408** sends, on behalf of Alice's device **3402**, a lookup request **3410** to Bob's secondary server **3412**. Bob's secondary server **3412** sends a lookup response **3414** to Alice's secondary server **3408**. Alice's secondary server **3408** forwards the lookup response **3414** as a forwarded response **3416** to Alice's device **3402**. Alice's device **3402** receives a copy of Bob's public key **3418** in the forwarded response **3416**.

(459) In step **1.4**, Alice's device **3402** encrypts the shared key **3405** with the copy of Bob's public key **3418** to generate an encrypted shared key **3420**.

(460) In step **1.5.1**, Alice's device **3402** initiates storing of the encrypted shared key **3420** on Alice's device **3402**.

(461) For instance, in step **1.5.2**, Alice's device **3402** sends an update request **3422** to Alice's secondary server **3408** to update @bob:key@alice with the encrypted shared key **3420**.

(462) In step **1.6.1**, Alice's device **3402** encrypts a phone@alice value Alice has shared with Bob (e.g., as @bob:phone@alice), using the shared key **3405**, to generate an encrypted phone value **3424**.

(463) In step **1.6.2**, Alice's device **3402** sends an update request **3426** to Alice's secondary server **3408** to update an @bob:phone@alice key with the encrypted phone value **3424**.

(464) At a later time, in step **2.1**, Bob wants to look up Alice's phone number.

(465) In step **2.2.1**, Bob's device **3404** initiates a lookup of the shared key shared between Bob and Alice.

(466) For example, in step **2.2.2**, Bob's device **3404** sends a lookup request **3428** to Bob's secondary server **3412** to lookup the shared key shared between Bob and Alice. Bob's secondary server **3412** forwards the lookup request **3428** to Alice's secondary server **3408**, as a forwarded request **3430**.

(467) In step **2.2.3**, Alice's secondary server **3408** sends an encrypted shared key to Bob's secondary server **3412**. Bob's secondary server forwards an encrypted shared key **3432** to Bob's device **3404**.

(468) In step **2.3**, Bob's device **3404** decrypts the received encrypted shared key **3432** using a private key stored on Bob's device **3404**, to obtain an unencrypted shared key.

(469) In step **2.4.1**, Bob's device **3404** sends a lookup request to Bob's secondary server **3412** to look up Alice's phone number that Alice has shared with Bob. Bob's secondary server **3412** can forward the lookup request to Alice's secondary server **3408**.

(470) In step 2.4.2, Alice's secondary server **3408** sends an encrypted phone value to Bob's secondary server **3412**. Bob's secondary server **3412** forwards an encrypted phone value **3434** to Bob's device **3404**.

(471) In step 2.5, Bob's device decrypts the encrypted phone value **3434**, using the shared key, to obtain the phone value for Alice.

(472) FIG. **35** illustrates an example system **3500** for key value retrieval.

(473) At a first stage, a user **3502** (e.g., user2) desires to obtain a key of a user **3504** (e.g., user1). A request for the key can be sent from a device **3506** of the user **3502** to a secondary server **3508** associated with the user **3504**.

(474) In a second stage, the user **3502** proves that s/he is user2 (e.g., using a pol approach as described herein).

(475) In a third stage, the secondary server **3508** locates an entry **3510** in a secondary storage **3512** that had been encrypted using an encryption key created by a settings application **3514** on a device **3516** of the user **3504**. The entry **3510** can include a key code (e.g., Key123), for example.

(476) In a fourth stage, the secondary server **3508** can send the entry **3510** (or a portion of the entry **3510**) with a request to the device **3516** for the settings application **3514** to decrypt the entry **3510**.

(477) In a fifth stage, the settings application **3514** can locate, e.g., in local storage **3518**, a key for the key code included in the entry **3510** and use the key to decrypt the entry **3510** (or to decrypt a value associated with the entry **3510**). The decrypted value is sent by the device **3516** to the secondary server **3508**. The secondary server can forward the value to the device **3506** so that the user **3502** can view the decrypted value (or otherwise use the decrypted value, e.g. in an application on the device **3506**).

(478) FIG. **36** illustrates approaches for spam protection. Various types of settings can be configured for controlling access to secondary servers. For example, some settings can relate to spam protection that can involve the use of an allow list **3602** and/or a block list **3604**. In a first example **3606** in which an allow list **3608** includes @alice and @bob and a block list **3610** is empty, connections can be accepted from @alice and @bob. In a second example **3612** in which an allow list **3614** includes a wildcard character '*' and a block list **3616** includes @daniel, connections can be accepted from everyone except @daniel. In a third example **3618** in which an allow list **3620** includes @alice and @bob and a block list **3622** includes @daniel, connections can be accepted from @alice and @bob but not @daniel. In a fourth example **3624** in which an allow list **3626** includes @alice and a block list **3628** includes connections from everyone except @alice can be rejected.

(479) In some implementations, an @config verb can support adding and removing entries from allow or block lists. For example, a verb "@config:block:add @bob @daniel" can result in an adding of @bob and @daniel to a block list for the user. As another example, a verb "@config:allow:remove:@bob" can remove @bob from an allow list of the user. Other examples can include "@config:show:block" and "@config:show:allow" to receive a list of blocked or allowed @signs, respectively.

Local and Secondary Server Synchronization

(480) FIG. **37** illustrates example synchronization operations **3700** for synchronizing local and cloud persistence. A settings application **3701** running on a user device **3702** can be involved in various synchronization operations.

(481) For example, in step **1**, the settings application **3701** can retrieve a last synced entry commit identifier from a local commit log **3704**.

(482) In step **2**, the settings application **3701** can retrieve a latest server commit identifier from a secondary server **3706**.

(483) In step **3**, the settings application **3701** can retrieve uncommitted entries from the local commit log **3704** based on the latest server commit identifier. As indicated in a note **3707**, the settings application **3701** can determine that a local key store **3708** is ahead of (e.g., has more

recent entries) than the secondary server **3706** if a local commit identifier is equal to a server commit identifier, the server commit identifier is a null value, or the local commit identifier is less than the server commit identifier and the number of uncommitted entries is greater than zero. As indicated by a note **3709**, the settings application **3701** can determine that the secondary server **3706** is ahead of the local key store if the server commit identifier is greater than the local commit identifier.

(484) In step **4**, the settings application **3701** can retrieve values for uncommitted keys from the local key store **3708**.

(485) Steps **5** and **6** can be repeated for each uncommitted entry. In step **5**, an update or delete command for a given uncommitted entry is sent by the settings application **3701** to the secondary server **3706**. The secondary server **3706** can return a current server commit identifier.

(486) In step **6**, the local commit log **3704** is updated for the given uncommitted entry to reflect the current server commit identifier.

(487) FIGS. **38A-C** illustrate a synchronization scenario **3800** in which a local storage **3802** is ahead of (update-wise) storage on a secondary server **3804**. In a first step, a value (e.g., value1) is inserted (e.g., as an entry **3806**) into the local storage **3802** for a key1 key. An entry **3808** for the key1 key can be added to a commit log (e.g., located on a same user device as the local storage **3802**). A secondary server time value **3810** for the commit log entry **3808** can be initially null, which can mean that the local storage entry **3806** needs to be committed to the secondary server **3804**. The entry **3808** can include an indication (e.g., illustrated as “+”) to represent an addition of a key value. Other indications can include an indication of deletion (e.g., “-”) or an indication of metadata change (e.g., “#”).

(488) In a second step, a local server that manages the local storage **3802** can determine a latest “optime” (e.g., operation time) of the secondary server **3804**. For instance, the local server can execute an @stat verb **3812**. A return value of null can indicate that the local storage **3802** is ahead of the storage on the secondary server **3804**.

(489) In a third step, the local server can run an update command **3814** to commit changes in the local storage **3802** to the secondary server **3804**. The update command **3814** can result in the secondary server **3804** adding an entry **3816** in e.g., cloud storage managed by the secondary server **3804**. The secondary server **3804** can also add an entry **3818** to a server commit log. The server commit log entry **3818** can include the key1 key, a timestamp T1, and an operation time (e.g., counter, having a value of 1).

(490) In a fourth step, the local server can update a cloud secondary time value **3820** in the commit log entry **3808**.

(491) In a fifth step, the local server performs additional operations that result in an addition of a key2 value (e.g., entry **3822**) and modification or addition of metadata for a key3 value (e.g., entry **3824**). Initially, secondary server time values **3826** and **3828** are empty for the key2 and key3 values, respectively.

(492) In a sixth step and as shown in FIG. **38B**, the local server performs an update command **3830** to commit the changes for the key2 key to the secondary server **3804**. The update command **3830** results in addition of an entry **3832** to the secondary storage and addition of an entry **3834** to the server commit log.

(493) The local server can run an @stat command **3836** to get a latest cloud secondary time (e.g., T2). Based on the output of the @stat command **3836**, the local server can update a cloud secondary time value **3838** for the key2 key in the local commit log. Similarly, the local server can perform an update command **3840** to commit the changes for the key3 key to the secondary server **3804**. The update command **3840** results in addition of an entry **3842** to the secondary storage and addition of an entry **3844** to the server commit log.

(494) The local server can run an @stat command **3846** to get a latest cloud secondary time (e.g., T3). Based on the output of the @stat command **3846**, the local server can update a cloud

secondary time value **3848** for the key3 key in the local commit log.

(495) In a seventh step and as shown in FIG. **38C**, the local server deletes a key (e.g. key1) from the local storage. An entry **3850** is added to the commit log to represent the deletion.

(496) In an eighth step, the local server executes a delete command **3852** to commit changes (e.g., including the key1 deletion) to the secondary server **3804**. An entry **3854** is added to a server commit log, and after another @stat command by the local server, a latest server optime value **3856** is added to the local commit log, to keep the local storage and the cloud storage in sync.

(497) FIG. **39** illustrates a synchronization scenario **3900** in which a local storage **3902** is behind (update wise) storage on a secondary server **3904**. In a first step, the secondary server has inserted two new values (e.g., for a key2 key and a key3 key, as shown in server entries **3906** and **3908**, respectively). Corresponding entries **3910** and **3912** have been added to a server commit log.

(498) In a second step, a local server executes an @stat command **3914** to determine a latest operation time (e.g., “optime”) of the secondary server. For instance, a T3 time corresponding to the entry **3912** is returned. The local server can compare the T3 latest server op time to a latest optime T1 **3916** that is currently stored in a local commit log. Based on the T3 value being greater than the T1 value, the local server can determine that the secondary server **3904** has more recent data than the local storage **3902**. In a third step, the local server executes a sync command **3918** that includes a reference to the T1 optime stored in the local commit log. The sync command **3918** requests data from the secondary server **3904** that has been stored on the secondary server **3904** since the T1 optime. Data returned from the sync command **3918** corresponds to data stored in the entries **3906** and **3908** on the secondary server **3904**.

(499) In a fourth step, the local server performs local update commands **3920** to add entries to the local storage **3902** for the key2 and key3 keys. Additionally, entries **3922** and **3924** are added to the local commit log, for the key2 and key3 keys, respectively. After the update commands **3920** are performed, the local storage **3902** is in sync with the secondary server **3904**.

(500) FIG. **40** illustrates a synchronization scenario **4000** in which a local storage **4002** becomes synchronized with a secondary server **4004** after data is deleted from the secondary server **4004**.

(501) In a first step, the local storage **4002** is in sync with the secondary server **4004**.

(502) In a second step, the secondary server **4004** deletes a key2 key and adds an entry **4006** to a server commit log. Accordingly, the local storage **4002** is behind (update wise) the secondary server **4004**.

(503) In a third step, a local server executes an @stat command **4008** to retrieve a latest secondary optime value (e.g., T4) from the secondary server **4004**. The local server can determine that the local storage **4002** is behind the secondary server **4004** based on the T4 latest server optime value being greater than a largest server optime value **4009** of T3 stored in a local commit log.

(504) In a fourth step, the local server performs a sync command **4010** to retrieve changes since the T3 optime value. The sync command **4010** results in the local server receiving key2 deletion information corresponding to the server commit log entry **4006**.

(505) In a fifth step, the local server executes a local update command **4012** that includes information received from execution of the sync command **4010**, to delete a key2 entry from local entries **4014**. The local server also adds an entry **4016** to the local commit log to reflect deletion of the key2 key from the local storage. The entry **4016** includes a T4 optime value reflecting a time of execution of the local update command **4012**.

(506) FIG. **41** illustrates a synchronization scenario **4100** in which a local storage **4102** becomes synchronized with a secondary server **4104** after data is deleted from the secondary server **4104** and added to the local storage **4102**.

(507) In a first step, the local storage **4102** is in sync with the secondary server **4104**.

(508) In a second step, the secondary server **4104** deletes a key2 key and adds a corresponding entry **4106** to a server commit log. A local server adds a key3 key to the local storage **4102** (e.g., an entry **4108**) and adds a corresponding entry **4110** to a local commit log.

(509) In a third step, a local server executes an @stat command **4112** to retrieve a latest secondary optime value (e.g., T3) from the secondary server **4104**. The local server can determine that the local storage **4102** is behind the secondary server **4104** based on the T3 latest server optime value being greater than a largest server optime value **4114** of T2 stored in the local commit log.

(510) In a fourth step, the local server performs a sync command **4116** to retrieve changes since the T2 optime value. The sync command **4116** results in the local server receiving key2 deletion information corresponding to the server commit log entry **4106**.

(511) In a fifth step, the local server executes a local delete command **4118** that includes information received from execution of the sync command **4116**, to delete a key2 entry from local entries **4119**. The local server also adds an entry **4120** to the local commit log to reflect deletion of the key2 key from the local storage **4102**. The entry **4116** includes a T4 optime value reflecting a time of execution of the local update command **4112**. The local server also sends an update command **4122** to the secondary server **4104** for the secondary server **4104** to add a key3 value to the secondary server **4104** (e.g., as shown by an entry **4124**). The secondary server **4104** also adds an entry **4126** to the server commit log to reflect addition of the key3 key. The local server updates an entry **4128** in the local commit log with a T5 value corresponding to a time of execution of the update command **4122**.

(512) FIGS. **42A-B** illustrate a synchronization scenario **4200** in which a local storage **4202** becomes synchronized with a secondary server **4204** after data is deleted from the local storage **4202** and added to the secondary server.

(513) In a first step, the local storage **4202** is in sync with the secondary server **4204**.

(514) In a second step, a local server deletes a key2 key from the local storage **4202** and adds a corresponding entry **4206** to a local commit log. The secondary server **4204** adds a key3 key (e.g., an entry **4208**) and a corresponding entry **4210** to a server commit log.

(515) In a third step, a local server executes an @stat command **4212** to retrieve a latest secondary optime value (e.g., T3) from the secondary server **4204**. The local server can determine that the local storage **4202** is behind the secondary server **4204** based on the T3 latest server optime value being greater than a largest server optime value **4214** of T2 stored in a local commit log.

(516) In a fourth step, the local server performs a sync command **4216** to retrieve changes on the secondary server **4204** since the T2 optime value. The sync command **4216** results in the local server receiving key3 addition information corresponding to the server commit log entry **4210**.

(517) In a fifth step, the local server executes a local update command **4218** that includes information received from execution of the sync command **4216**, to add a key3 entry to the local storage **4202**. The local server also adds a corresponding entry **4222** to the local commit log to reflect addition of the key3 key to the local storage **4202**.

(518) In a sixth step, the local server executes a local update command **4224** to update a commit log entry **4226** for the key3 key. The local server **4202** sends a delete command **4228** to the secondary server **4204** for the secondary server **4204** to delete the key2 key. The secondary server deletes the key2 key and adds a corresponding entry **4230** to the server commit log. The local server sends an update command **4232** to the secondary server **4204** for the secondary server **4204** to add a server commit log entry **4234** for the key3 key, to reflect a time of adding the key3 key to the local storage **4202**. The local server updates a local commit log entry **4236** to reflect a time of deletion of the key2 key.

(519) FIG. **43** illustrates an example system **4300** for synchronizing keys across user devices and a secondary server **4302**. For instance, a user **4304** has a first user device **4306** and a second user device **4308**. The first user device **4306** has a local storage **4310** that includes an entry **4312** for a key123 key. The second user device **4308** has a local storage **4314** that includes an entry **4316** for a key456 key. The user can use a settings application **4318** on the first user device **4306** and/or a settings application **4320** on the second user device **4308** to execute a sync keys command. Execution of the sync keys command(s) can result in an added entry **4322** in the local storage **4310**

for the key456 key and an added entry **4324** in the local storage **4314** for the key123 key. In some implementations, synchronization can involve using data from an entry **4326** and/or an entry **4328** in secondary storage **4330** on the secondary server **4302**.

(520) FIGS. **44A-B** illustrate an example system **4400** for synchronizing keys across applications and a secondary server **4402**. In a first stage, a first application **4404** (e.g., an @me application) has added entries **4406** and **4408** for phone.me@alice and location.me@alice keys, respectively, to local storage **4410**. Corresponding entries **4412** and **4414** have been added to a local commit log. A server commit log **4416** is currently empty.

(521) In a second stage, a local server executes an @stat command **4418** to get a latest server optime. In this example, a null value is returned. Accordingly, the local server can determine that the local storage **4410** has more recent data than the secondary server **4402**. The local server can send an update command **4420** to the secondary server **4402** for the secondary server **4402** to add the phone.me@alice key. Accordingly, the secondary server **4402** has added an entry **4422**. The local server can update the local commit log entry **4412** to include an optime of one. The local server can send an update command **4426** to the secondary server **4402** for the secondary server **4402** to add the location.me@alice key. Accordingly, the secondary server **4402** has added an entry **4428**. The local server can update the local commit log entry **4414** to include an optime of two.

(522) In a third stage and as shown in FIG. **44B**, a second application **4432** (e.g., an @buzz application) has added an entry **4434** to a local storage **4436** and a corresponding entry **4438** to a local commit log. The local storage **4436** can be a different portion of a same storage area as the local storage **4410**, or can be a different storage repository. As another example, the second application **4432** can be running on a different device than the first application **4404** and the local storage **4410** and the local storage **4436** can reside on respective devices. Accordingly, each respective device can run a different, respective local server. If the two applications are running on the same device, each application can have a different local server instance. Additionally, each application can have a different local commit log (or can access different portions of a same local commit log).

(523) In a fourth stage, the local server executes an @stat command **4440** to obtain a latest server optime. Based on a results on the @stat command **4440**, the local server can determine that the secondary server **4402** has more recent data than the local storage **4436**. Accordingly, the local server can send a sync command **4442** to the secondary server **4402** to obtain a data update from the secondary server **4402**. As indicated by a note **4444**, the local server can determine that data returned from the sync command **4442** does not include any keys used by or associated with the second application **4432**. Accordingly, the local server does not update the local commit log or the local storage **4436**. The local server can send an update command **4446** to the secondary server **4402** for the secondary server **4402** to add an email.buzz@alice key based on information in the entry **4434**. Accordingly, the secondary server **4402** has added an entry to secondary storage for the email.buzz@alice key and a corresponding entry **4448** to the server commit log. The local server can add an optime value **4450** the local commit log entry **4438**.

(524) FIG. **45** illustrates an example sync response message **4500**. The sync response message **4500** includes a first portion **4502** for a first operation **4504** and a second portion **4506** for a second operation **4508**. The first operation **4504** corresponds to a key **4512** having a value **4514**. The second operation **4508** corresponds to a key **4516** having a value **4518**. An expires-at value **4520** (e.g., a TTL value) can indicate that the value expires at a specified time point. An available-at value **4522** (e.g., TTB value) can indicate that the value **4518** becomes available at a certain time point.

Design

(525) FIG. **46** illustrates example entities for verb handling. A secondary server entity **4602** can be associated with a verb handler manager **4604** and a verb executor **4606**. The verb handler manager **4604** can manage one or more verb handler entities **4608**. A verb handler entity **4608** can handle a

verb **4610**. A verb **4610** can be received at the secondary server entity **4602** using an inbound connection **4612**. A verb can be executed at the secondary server entity **4602** by the verb executor **4606**. The verb handler entity **4608** can store output from verb execution, e.g., as may be produced by some verbs, in a secondary store **4614**. As another example, the verb handler entity **4608** can retrieve information from the secondary store **4614**, e.g., for use during verb execution. When execution of the verb results in output to be provided to a requestor, the verb handler entity **4608** can provide generated output to the requester, e.g., using an outbound connection **4616**.

(526) FIG. **47** illustrates an example class diagram **4700**. A key store manager interface **4702** can be based on a common interface **4704** that represents commonality between a read only key store interface **4706** and a writable key store interface **4708**. A common library **4710** (used by some or all depicted interfaces) can include key store operations **4712** and exception handling functionality **4714**.

(527) FIG. **48** illustrates client design criteria goals **4800**. A first goal **4802** specifies that a developer of a protocol application need not be exposed to internal @protocol internals or terminology. For instance, a lookup interface in an @protocol API may be designed to be more simple than a get interface included in other API(s). More precisely, a goal **4804** can specify an approach of designing for simple method signatures. Signature and other method considerations can include use (or avoidance) of exception handling, return values, and callbacks. A third goal **4806** specifies that developers be aware of what interface may result in execution on an @user's secondary server or on a user's device. A fourth goal **4808** specifies that designing for widget (UI) usability is desired.

Notifications

(528) The @protocol can provide a method for secure, surveillance-free, P2P notifications with privacy control features. Other notifications that may be sent over the Internet generally require an intermediate service to manage queueing, routing, and delivery. Such existing services can rely on possession of unique endpoint identifiers (e.g., phone numbers) to determine routing. Such factors can result in surveillance and tracking of notifications which is not desirable. In contrast, the @protocol includes a mechanism for pushing notifications between @sign servers the instant a data change event has occurred for a data item of interest. These automatic notifications can provide an essential element of creating a P2P application model with data distributed amongst @servers. The @protocol can also provide a general purpose notification service that application developers can use for application functionality to securely and safely push notifications to other @signs as part of application behavior.

(529) The @protocol is designed to give data owners control over who can access their data. With the @protocol, a requester's identity is verified before data of a data owner is released to the requester. The @protocol is thus different from existing protocols which do not have such verification. With other protocols, unwanted communication (e.g., spam) is a continuing concern and disadvantage.

(530) The @protocol can manage notifications regarding updates to information. Particularly, the @protocol distributes updates differently than other protocols and existing techniques. For example, consider when a user, Alice, changes her location. Although a user Bob can look up Alice's location through a previously known @protocol address, such as location@alice, Bob may not know if Alice's location has changed unless he looks it up again. However, it can be unreasonable to expect Bob to constantly be checking Alice's location on the off chance that Alice's location may change. As an alternative approach, push notifications can be used.

(531) Push notifications can allow Alice to alert Bob when her location or other Alice-owned data has changed. Currently, most organizations update customers by pushing out emails, phone calls, and SMS, which may result in spam. Not only may spam be annoying and time-wasting for recipients, spam also (unfairly) consumes receiver's resources. A more fair approach may be to have a sender bear the cost of sending updates (e.g., similar to real-world postal mail). Accordingly,

the @protocol can provide real time notifications of updates with the sender/data owner generally incurring a cost of sending and storage. With the @protocol, users only have to update information in one place.

(532) The @protocol can use an approach that may be similar to raising a flag on a receiver's mailbox as an indication of updated data. That is, @protocol notifications can involve sending notification flags to other @signs. The receiver can decide whether to respond to the notification. If the receiving @sign is interested in the notification, the receiving @sign can choose to review the details of that notification. As described below, the @protocol can handle notifications for both online and offline devices.

(533) Various considerations and goals can apply to notifications. For example, if Alice wants to connect to Bob, then she can notify Bob of that intent using the @protocol, and Bob can be notified of that intent and be able to respond to the notification. As another example, Bob may or may not want to respond to requests from @signs he does not have a relationship with, or maybe Bob only wants to respond to @signs to which he had given permission. As yet another example, some users may desire to automatically agree to connections, but using a particular persona (e.g., a public persona) in those instances, thereby providing some basic information without providing more private or otherwise non-public information. In general, the @protocol is designed to avoid and protect from mass spamming of requests from @signs. Some notifications, such as request notifications, can be used after a "connection" between two parties has been established. Request notifications can have a particular context. For example, Alice may want to send a message to Bob, and a context in this case is a need to send a message. A response can be or include a response of accept, reject, ignore, or further context in reply. For example, a reply with further context can be, "I can't message now, but call me in 10 minutes using messaging application ABC."

(534) FIG. 49 illustrates an example system for an accept notification. A user @alice **4902** wants to connect with a user @bob **4904**. A device **4906** of the @alice user **4902** can send an invitation to connect notification **4908** to the @bob user **4904**. The invitation to connect notification **4908** can be provided to a device **4910** of the @bob user **4904**. As described below, Alice's secondary server and/or Bob's secondary server can be involved in distribution of notifications. The @bob user **4904** can choose whether to accept and/or respond to the invitation to connect notification **4908**. If the @bob user **4904** accepts the invitation to connect notification **4908**, the device **4910** can send an accept notification **4912**, which can be provided to the device **4906** for viewing by the @alice user **4902**.

(535) FIG. 50 illustrates an example system **5000** for an update notification. A user Alice (@alice **5002**) may update a key, such as phone@alice. The @alice user **5002** may or may not want to notify other users of that change, or may want to notify only certain users. Alice may choose to send individual notifications to particular users, or may instead send group- or persona-based notifications. For instance, Alice may get a new phone number and may choose to notify groups of friend or work colleague users, or users to which a friend or work persona has been shared. In this example, Alice may choose to notify an @bob user **5004** of the new phone number.

(536) A device **5006** of the @alice user **5002** can send an update notification **5008** to the @bob user **5004** regarding a change to the phone@alice key. The update notification **5008** can be provided to a device **5010** of the @bob user **5004**. The @bob user **5004** can choose whether to respond to the update notification **5008**. If the @bob user **5004** chooses to respond to the update notification **5008**, the device **5010** can send a response message **5012**, which can be provided to the device **5006** for viewing by the @alice user **5002**.

(537) FIG. 51 illustrates a peer-to-peer network **5100** for group notifications. Some notifications can be sent to all members of a group. Devices **5102**, **5104**, **5106**, **5108**, and **5110** of group users can be configured to be in a peer-to-peer relationship, where each device in the group is trusted by the other devices and can therefore be used to send messages to the devices in the group. A consideration for a peer-to-peer network is maintenance of group member lists by each device in

the network. An alternative approach for group management is a reflector architecture.

(538) FIG. 52 illustrates an example reflector network 5200. The reflector network 5200 can include a reflector device 5202. The reflector device 5202 corresponds to an entity that is trusted by everyone in the group. Group members (corresponding to devices 5204, 5206, 5208, and 5210) do not need to know or have any relationship to each other, beyond each trusting the same reflector. The @protocol can use a reflector architecture, where a group owner adds members to the group and everyone in the group can receive notifications from the reflector device 5202.

(539) FIG. 53 illustrates information 5300 describing a notification verb 5302. The notification verb 5302 can be used to notify another user's secondary server of some activity (e.g., the notification verb can be provided by the @protocol as a method for secondary servers to notify other secondary servers about certain events). A notification can turn into some kind of action by the recipient.

(540) A syntax 5304 of the notification verb can be:

(541) notify:send<whom_to_notify>:<key>@<owner>.

(542) For instance, an example notification verb can be:

(543) notify:send@kevin:phone@colin.

(544) An overall flow 5306 of notification processing can include specification of whom to notify, what to notify about, enabling an action, or just sending to inform, where action(s) in response to a notification can be automated or manual.

(545) Example notification scenarios 5308 include an @bob user sharing a persona with @alice and notifying @alice when a value in the persona has changed. Other examples include the @bob user can notifying one or more users of a new chat message or a new email message.

(546) Notification verb features 5310 include that notifications should not spam recipient(s), notifications can be streamed (e.g., to application(s) to enable application reaction(s), notifications can have an expiration, automatic notifications can be supported, and a notification can be represented as a reference (e.g., using an atsign://link).

(547) FIG. 54 illustrates an example system 5400 for notifications. Rather than simple polling or push notifications schemes, the @protocol can use secondary servers to capture notifications from other @signs in real time and create a queue of notifications on a sender's secondary server.

Notifications can be polymorphic in that acted-on notifications can refer to a same or different locations in a notifier's namespace. In the example of FIG. 54, a user Bob 5402 (@bob) wants to let a user Alice 5404 (@alice) know he has a new phone number listed at phone@bob.

(548) In a first stage 5406, the user Bob 5402 authenticates to an @bob secondary server 5408. The user Bob 5402 makes an update to his phone@bob key (e.g., by sending an update command to the @bob secondary server 5408 from a device 5410) to reflect his new phone number. The user Bob 5402 then uses the device 5410 to send a notification verb to @alice.

(549) In a second stage 5412, the @bob secondary server 5408 connects to an @alice secondary server 5414 on behalf of @bob, authenticates as @bob to the @alice secondary server 5414 (e.g., using from and pol verbs), and then submits a notification verb that includes phone@bob to notify @alice of an update to phone@bob.

(550) In a third stage 5416, a device 5418 for the @alice user 5404 synchronizes to the @alice secondary server 5414. As part of synchronization, the notification for phone@bob can be provided to the device 5418. If the notification is related to a root namespace, a settings application on the device 5418 can automatically handle the notification. As another example, if the notification is in a registered application namespace, an application associated with the namespace can handle the notification. Other examples include a notification message being displayed on the device 5418.

(551) In a fourth stage 5420, a notification flag can be cleared, for example, in response to a time out or in response to the @alice secondary server 5414 looking up the phone@bob key.

(552) FIG. 55 illustrates an example system 5500 for notifications. In an example scenario, a user Bob 5502 (@bob) wants to notify a user Alice 5504 (@alice) of a new message. The user Bob 5502

can update a key message@bob and notify the user Alice 5504. A device 5506 of the Alice user 5504 can be powered on and monitoring for real time updates.

(553) In further detail, in step 1, the device 5506 for Alice authenticates two sessions to an @alice secondary server 5508 using a pkam verb.

(554) In step 2 and in the first session, the device 5506 issues a monitor verb to the @alice secondary server 5508. The monitor verb can be submitted without arguments. In general, once a monitor verb is invoked, a connection/channel can be established with a secondary server and notifications can be sent directly over the connection. Notifications on the channel can be treated as events that can be acted on immediately by an application or ignored. After the device 5506 invokes the monitor verb, a stream of real-time notifications can be sent to the device 5506. The format of a notification sent by a monitor can be “data:<service>@<entity>” (e.g., data:message@bob).

(555) In step 3 and in the second session, the device 5506 can perform a lookup on the @alice secondary server 5508 for a q.alice@alice key to determine whether there are any notifications that have been received by the @alice secondary server 5508 while the device 5506 was offline. For example, if a notification of data:message@bob had been received by the secondary server 5508 when the device 5506 was offline, the lookup:q.alice@alice command can return data:message@bob. However, if multiple notifications are in the queue for @alice, then the lookup:q.alice@alice command can return a list:reply, such as list:5 message@bob, which can indicate that the queue holds five messages and a first value is message@bob. A next invocation of the lookup command may return, for example, list:4 message@sara, for example, with a second item in the queue being a notification about a message from @sara. Once a list is exhausted then a lookup command can return data:null.

(556) In step 4, the device 5506 is ready and listening for real time notifications. As an example, the user Bob 5502 has a notification to send corresponding to an update to a key @alice:message@bob. For instance, the user Bob 5502 may want to notify the user Alice 5504 of a message targeted for Alice. Bob can send notifications manually, or, in some implementations, notifications can be sent automatically (e.g., after a config:notify: auto configuration). Before sending the notification, a device 5510 for @bob authenticates to an @bob secondary server 5512 using the pkam verb.

(557) In step 5, the device 5510 issues a notify:@alice:message@bob verb.

(558) In step 6, the @bob secondary server 5512 authenticates to the @alice secondary server 5508 using the from and pol verbs. In step 7, the @bob secondary server 5512 issues a notify:message@bob verb to the @alice secondary server 5508. In step 8, based on the monitoring channel with the device 5506, a real time update of data:message@bob is made using the monitoring connection from the @alice secondary server 5508 to the device 5506.

(559) In step 9, and in response to the real time update, the device 5506 triggers an event for an @protocol application on the device 5506. The application is interested in (e.g., has subscribed to) the notification, and can, for example, store the notification for later processing (e.g., for when Alice 5504 may want to acknowledge the notification and have the application process the event). As another example, the application can perform a lookup (e.g., looking up a value for message@bob that is shared by @bob with @alice) so that the contents of the message are ready for presentation to Alice when Alice next uses the application.

(560) In step 10, either the application on the device 5506 or the Alice user 5504 wants to know more about the notification. The device authenticates to the @alice secondary server 5508 using the pkam verb. In step 11, the device 5506 issues a lookup:message@bob request to the @alice secondary server 5508. In step 12, the @alice secondary server 5508 authenticates to the @bob secondary server 5512 using the from and pol verbs. In step 13, the @alice secondary server 5508 issues a lookup:message@bob verb to the @bob secondary server 5512. In step 14, the @bob secondary server 5512 looks up the message@bob key shared with @alice and sends a looked up

value (e.g., “Hi Alice”) to the @alice secondary server **5508**, e.g., as “data:Hi Alice.” As another example, if the message@bob value shared with @alice is a list, a list indicator (and first list item), such as “list:10 Hi Alice” can be provided. The @alice secondary server **5508** can forward the received value for the message@bob key to the device **5506**, for use by or presentation in an application on the device **5506**. If a list is returned, the @alice secondary server can cycle through the items in the list and provide each list item to the device **5506**.

Example Applications

(561) An ability of the @protocol to locate edge resource endpoints that securely store a user's data gives rise to a new class of applications that can communicate with each other directly (e.g., P2P). This is an important capability in that it removes a scenario for some intermediary (e.g., typically called a server, but could also be a more nefarious man in the middle scenario) being used to store data and manage interactions with other users that want to access stored data. Additionally, end-to-end encryption ensures that all data and traffic only flows between the involved parties and is not susceptible to surveillance. Accordingly, the @protocol can support a new generation of true P2P applications that can securely share information between permitted users without fear of surveillance by intermediaries, infrastructure providers, or others. Additionally, registering for a new protocol-enabled application can be a simple process for users, because a user can register using their @sign which can reduce sign-up/registration steps and can enable user-controlled access of data to application providers/developers, rather than a user having to separately and repeatedly provide multiple specific items of user information each time a user installs or registers for an application.

(562) FIG. **56** illustrates an example system **5600** for P2P applications. An application **5602** running on a device **5604** of a user @bob can request data for an @alice user stored, e.g., in a personal directory **5606** of a secondary server **5608** and/or a personal directory **5610** on a device **5612** of @alice. Similarly, an application **5614** running on the device **5612** of @alice can request data for the @bob user stored, e.g., in a personal directory **5616** of a secondary server **5618** and/or a personal directory **5620** on the device **5604** of @bob. Specific application examples are described in more detail below.

(563) Data sharing between the device **5604** and the device **5612** (e.g., and/or between the application **5614** and the application **5602**) can be conducted without a need of or actual surveillance by a central server. P2P application functionality can be configured with no central server or entity being able to track a user's use of or sharing of data. Data that a user or user application determines to share with another user or application is encrypted, as described above. Also as described above, one client application of one @protocol user can be notified about a change in data for another @protocol user. As described in more detail below, an example P2P application can be a contacts application, in which each user is responsible for maintaining their own data, and interested parties/applications can be automatically notified about changes in, for example, subscribed-to data. Another example application is the file sharing application described above. Yet another example application is a data sharing application for use by an executive and an executive assistant.

(564) FIGS. **57A-B** are conceptual diagrams **5700** and **5701** that illustrate an example P2P application for sharing data between an executive and an executive assistant. As shown in FIG. **57A**, an executive **5702** may occasionally have a need to share data with an executive assistant (EA) **5704**. For example, the executive **5702** may have various sensitive items of information, such as a social security number **5706**, travel information **5708**, family information **5710**, passwords **5712**, or credit card information **5712**. The EA **5704** may occasionally need, when supporting the executive **5702**, a sensitive item of information of the executive **5702**, such as for making a reservation for the executive **5702** or performing some other type of transaction on behalf of the executive **5702**, such as filling out a visa application. Prior, insecure solutions for sharing sensitive information between an executive and an EA may include using paper file folders, notes, or other

types of paper exchanges, e-mail messages or vCards (virtual cards), which could inadvertently or maliciously be forwarded, or an electronic file, which could also be shared to or accessed by unauthorized individual. These prior approaches are not sufficient in terms of securing sensitive information of the executive **5702**.

(565) As an alternative for sharing data between the executive **5702** and the EA **5704**, an @protocol-enabled P2P data sharing application can be used where the executive **5702** owns and controls her data and can give access to select individuals, such as the EA **5704**, for a period of time and/or on an as-needed basis. Data requests and accesses can be logged, thereby enabling review by the executive **5702**.

(566) The data sharing application can have an executive mode or variant for running on an executive device **5716**, and an EA mode or variant for running on an EA device **5718**. The EA **5704** can use the EA application variant on the EA device **5718** to generate a request for sensitive information of the executive **5702** (e.g. a request **5722** for a credit card number). As shown in FIG. **57B**, in response to a validation of the request **5722**, requested data **5724** can be provided to the EA device **5718** and presented on the EA device **5718**, e.g., for a limited period of time. The @protocol-enabled P2P data sharing application is further described below.

(567) FIG. **58** illustrates an example configuration user interface **5800** for a P2P data sharing application. The configuration user interface **5800** can be presented upon first use of the P2P data sharing application, for example. The user can perform a role selection by selecting either an executive option **5802** or an EA option **5804**.

(568) FIG. **59** illustrates an example executive assistant list user interface for a P2P data sharing application. The executive assistant list user interface **5900**, which can be available for an executive user of the P2P data sharing application with an executive role, presents a list of EAs that may support the executive user. For example, the list of EAs includes Larry Martin **5902**, June Siegel **5904**, and Jen Waterly **5906**. The executive can select a particular EA to view information about the EA and/or to configure access for the EA.

(569) FIG. **60** illustrates an example executive assistant permissions user interface **6000** for a P2P data sharing application. The executive assistant permissions user interface **6000** presents permissions that have been (or that can be) granted to a selected executive assistant by an executive user of the P2P data sharing application. For example, data that can be shared with the Jen Waterly EA includes a social security number **6002**, Visa credit card information **6004**, American Express credit card information **6006**, and travel site login information **6008**. Toggle controls **6010**, **6012**, **6014**, and **6016** can be used to toggle access to the social security number **6002**, the Visa credit card information **6004**, the American Express credit card information **6006**, and travel site login information **6008**, respectively.

(570) FIG. **61** illustrates an example executive list user interface **6100** for a P2P data sharing application. The executive list user interface **6100**, which can be available for an EA user of the P2P data sharing application with an EA role, presents a list of executives that the EA supports. For example, the list of executives includes Jan Lorel **6102**, Mark Wayfair **6104**, and Lauren James **6106**. The EA can select a particular executive to view information about the executive and/or to view information about items that the executive has shared with the EA or to request access to a certain item of information.

(571) FIG. **62** illustrates an example data access user interface **6200** for a P2P data sharing application. The data access user interface **6200**, which can be available for an EA user of the P2P data sharing application with an EA role, presents a list of shared data items that a selected executive has shared with the EA. For example, the list of shared data items includes a shared social security number **6202**, shared Visa credit card information **6204**, shared American Express credit card information **6206**, and shared travel site login information **6208**. The EA can select a particular shared item to view a shared data item. In some implementations, a shared data item is masked so that the EA cannot view the actual value. The data sharing application may enable the

value to be copied (e.g. into another application, which may also mask the value) but not viewed, for example. As mentioned, data access can be logged. In some implementations or for some data items or data sharing, a shared data item may be accessible for a limited period of time.

(572) FIG. **63** is a block diagram illustrating an example system **6300** for execution of a P2P data sharing application. An executive data sharing application **6302** can be installed, with an executive role, on an executive mobile device **6304**, e.g., for an executive protocol user @lauren. The @lauren executive user can be a registered @protocol user, and the mobile device **6304** can include a personal directory **6306** that is synchronized with a personal directory **6308** on a secondary server **6310** configured for @lauren. The personal directories **6306** and **6308** can include shared travel login information **6312** or **6314**, respectively, that has been shared, for example, through the executive data sharing application **6302** with an EA user (e.g., @larry) associated with an EA mobile device **6316**. The EA mobile device **6316** includes an executive data sharing application **6318** running with an EA role.

(573) The @protocol can enable the @lauren executive user to share, using the executive data sharing application **6302**, sensitive information, such as the shared travel login information **6308**, with the @lauren user. The @protocol can enable the @lauren user to access the shared data using the executive data sharing application **6318**. The @lauren user can also be a registered @protocol user, with a configured secondary server **6320** and a personal directory **6322** (that can be synchronized with a personal directory **6324** on the EA mobile device **6316**).

(574) The executive data sharing application **6318** can use the @protocol to send a request, over a network **6326**, to the secondary server **6320**, for the secondary server **6320** to request an item of information, such as the shared travel login information **6308**, that the @lauren user has shared with the @larry user. The secondary server **6320** for @larry can send a lookup request to the secondary server **6310** for @lauren. The secondary server **6310** for @lauren can validate the request, and if the requested executive data has indeed been shared with the EA @larry user, the secondary server **6310** can provide the shared information to the secondary server **6320**, for forwarding on to the executive data sharing application **6318**. The @larry user can view the shared data in the executive data sharing application **6318**, and/or use the shared data in another application **6328** (e.g., a travel reservation application).

(575) FIG. **64** illustrates a user interface **6400** for an improved contacts application. With traditional contacts applications, data stored by the contact application generally becomes progressively out of date. For instance, a first user may receive contact information from another user, but a responsibility for maintaining the contact information on a device of the first user generally falls to the first user. This can be problematic because the first user may not know that the other user's contact information has changed. The problem can be compounded because the first user may have contact information for many other users, with each set of contact information potentially (or eventually) becoming out of date. Additionally, when the first user's own contact information changes, the first user may not remember who they have shared their contact information with, so may not know who to provide updated contact information.

(576) To solve these and other problems, users can use the improved contacts application, which can be a surveillance free P2P protocol-enabled application used by each user. Each user can choose to whom to share contact and other information. Each user can automatically view current contact information from other users, each time contact information is accessed.

(577) The user interface **6400** presents information for an @chloe user **6402**. The improved contacts information, which may have a name of “@buzz,” uses a theme based on bees. A waggle dance area **6404** shows invitation-related messages, such as invitations **6406** to connect or share from other protocol users, event invitations **6408**, or other invitations. A swarming area **6410** corresponds to groups or topics to which the @chloe user **6402** belongs or subscribes. For example, the @chloe user **6402** is a member of a stats study group **6412**, is interested in a set of topics **6414**, and participates in e.g., a group chat **6416**. A reminders area **6424** provides indications of contacts

the user **6402** has not recently contacted.

(578) The user interface **6400** may be a home screen for the improved contacts application. The user **6402** can navigate to the home screen by selecting a “for you” item **6426** in a navigation pane **6428**. A user can view more information about relationships with other protocol users by selecting a relationships item **6430**. Relationships with groups of connected users can be explored by selecting a hives item **6432**. Additional information regarding groups to which the user **6402** belongs can be explored by selecting a swarms item **6434**. A settings interface for the improved contacts application can be displayed in response to selection of a settings item **6436**. In some implementations, a bee **6438** can be selected as a shortcut for a particular operation. For instance, the bee **6438** can be selected to initiate creation of a new swarm (group) of users.

(579) FIG. **65** illustrates an example invitation user interface **6500** for the improved contacts application. The invitation user interface **6500** can be displayed in response to selection of invitations **6406** on the user interface **6400** described above with respect to FIG. **64**. The invitations user interface **6500** is displaying invitations **6502** from other protocol users (e.g., as part of a people area **6504**). For example, the invitations **6502** include an invitation **6505** from an @naomi user. Invitations from organizations or groups can be displayed in response to selection of an organizations item **6508** or a hives item **6510**, respectively.

(580) FIG. **66** illustrates an example relationship user interface **6600** for displaying relationship information for the improved contacts application. The relationship user interface **6600** can be displayed in response to selection of a relationships item **6602** in a navigation pane **6604**. An invitation area **6606** displays connection or sharing invitations, such as an invitation **6608** from the @naomi user. A reminders area **6610** provides indications of contacts the user has not recently contacted. A suggestions area **6612** displays suggestions for connecting to other protocol users (e.g., based on common connections). The user may respond to a particular invitation in the invitations area **6606** by selecting the invitation. For example, may select the invitation **6608** to respond to an invitation from @naomi.

(581) FIG. **67** illustrates an example persona selection user interface **6700**. The persona selection user interface **6700** can be displayed in response to accepting an invitation to connect with another protocol user (e.g., the @naomi user). The persona selection user interface **6700** enables the user to select, from among a set of personas the user has configured, a particular persona to share with the requesting user. A persona can be a group of items or sharing settings that the user wishes to define for sharing as a group of items (e.g., a particular phone number, whether to share location, whether to share social media activity, etc.). The user can share a particular persona with multiple connections that the user views similarly. For example, the user has defined a Ruby's human persona **6702** (e.g., rubys-human.profile@chloe), a BFF (Best Friends Forever) persona **6704** (e.g., bff.profile@chloe), a social media persona **6706** (e.g., social.profile@chloe), a family persona **6708** (e.g., family.profile@chloe), and a school persona **6710** (e.g., schoo.profile@chloe), for sharing, e.g., settings for dog park acquaintances, close friends, other social media acquaintances, family members, and school friends, respectively. For connecting with @naomi, the user can select one of the configured personas to share with @naomi. For instance, based on meeting @naomi at a dog park, the user can select the Ruby's human persona **6702**.

(582) FIG. **68** illustrates an example user interface for viewing personas. As described above, a user may define different personas for sharing different sets of information or establishing different sharing settings, for different situations or for different connected users or sets of users. The user has selected a Ruby's human persona **6802**. Other personas include a BFF persona **6804**, a social media persona **6806**, a family persona **6808**, a school persona **6810**, a Linked-in persona **6812**, an Instagram persona **6814**, and a Facebook persona **6816**. For the selected Ruby's persona **6802**, an interaction count **6818**, a relationship count **6820**, and a hives count **6822** are displayed. The persona **6802** can include or be associated with an image **6824** and a description **6826**. Further information and settings for the persona **6802** can be viewed by selecting a control **6828**.

(583) FIG. 69 illustrates an example persona user interface **6900**. The user can use the persona user interface **6900** to configure a persona **6902**. For example, a control **6904** can be selected to configure a persona profile picture **6906**. As another example, a control **6908** can be selected to configure a persona description **6910**. The user can configure other settings that can control what information is shared with users to which the persona **6902** is shared. For example, a calls setting **6912** can be used to configure whether the user will accept calls from users to which the persona is shared. Social media **6914**, messages **6916**, and photo/video **6918** settings can be selected to configure whether social media information, messages, or photo/video content is shared with users to which the persona is shared, respectively. As another example, a location setting **6920** can be configured that can control whether the user's location is shared with users to which the persona has been shared and/or whether those users are notified when the user is at a same location. In addition to the displayed settings, the user can configure whether respective data items, such as a phone number or address, are shared when the persona is shared. In summary, the persona **6902** represents a group of settings that can be shared, as a group, to various protocol users or entities. When the user wants to share a same set of information to different users, the persona **6902** can be a convenient, efficient approach.

(584) FIG. 70 illustrates an example relationships user interface **7000**. The relationships user interface **7000** can display, for example, information regarding which personas have been shared with which protocol users. For instance, a BFFs area **7002** indicates users **7004** to which a BFF persona **7006** has been shared. Sharing the BFF persona **7006** with the users **7004** can result in sharing of photos **7008**, videos **7010**, posts **7012**, and messages **7014**, for instance. As another example, a family area **7016** indicates users **7018** to which a family persona **7020** has been shared. Sharing the family persona **7020** with the users **7018** can result in sharing of messages **7022**, mail **7024**, and calendar data **7026**, for example.

(585) FIG. 71 illustrates an example group creation user interface **7100**. A nearby users area **7102** displays protocol users who have been detected to be within a proximity of the user. The nearby users area includes users **7104**, **7106**, **7108**, **7110**, **7112**, **7114**, and **7116**, for example. The user can select which users to include in a new group. For instance, the user has selected the users **7104**, **7106**, **7108**, **7110**, **7112**, and **7114**.

(586) FIG. 72 illustrates an example group creation user interface **7200**. The user interface **7200** can be used to create a group (e.g., swarm) of users based on proximity, for example. An item **7202** can be selected to send a notification (e.g., broadcast message, using, e.g., low power Bluetooth) to nearby users, as an invitation to join a new group. Items **7204** or **7206** can be selected to send an invitation using a QR (Quick Response) code or short code, respectively. An item **7208** can be selected to select which person to share with the new group.

(587) FIG. 73 illustrates an example group creation user interface. A group users area **7302** displays users that are included in a new group. The group users area **7302** includes users **7304**, **7306**, **7308**, **7310**, **7312**, and **7314**, for example. The user can select which persona is shared with the users of the new group.

(588) FIG. 74 illustrates an example system **7400** for attestations. The @protocol can support trusted attestations. Attestations are evidence or proof of a claim made by a claimant (e.g., a claimant that is an @alice @protocol user with a mobile device **7402** and a secondary server **7404**). The claimant may provide claims such as of a pattern of "claim1@alice" or "claim2@alice" as attestations. A self-attestation can return a response by value. For example, a requester (e.g., a requester @protocol user who may be e.g., @bob, with a client device **7406** and a secondary server **7408**) may submit a lookup **7409** of an attestation such as claim1@alice. The secondary server **7404** for @alice can provide a response value **7410**, which can be provided, e.g., to the secondary server **7408** and then to the client device **7406**. The response value **74110** can be polymorphically determined, e.g., having a value determined based on an identity of the requester.

(589) An attestation may be deemed to be more trustworthy when a value for the attestation is

provided by a reliable third party, such as a trusted entity **7412**. The trusted entity **7412** can provide a value or can provide proof of some claim (e.g., an attestation) without the need for exposing the underlying data that the claim relies upon. For example, The U.S. Department of State might attest that @alice holds a valid passport without having to divulge any further information other than that the passport for @alice exists and is valid. In general, a trusted service **7414** provided by the trusted entity **7412** can validate a claim by any appropriate means available to the trusted entity **7414**. The appropriate means can include whatever appropriate tests or validations the trusted entity **7414** deems as necessary to validate the attestation. Example attestations that may be serviced by the trusted entity **7412** (or another trusted entity) can include attestations for which the trusted entity **7412** returns a true or false answer, such as ishuman@alice, isAdult@alice, hasSSN@alice, hasDriverLicense@alice, etc.

(590) The @protocol can enable a protocol user to create a trusted entity attestation using reference value(s). For example, in response to a lookup **7416** for claim2@alice, the secondary server **7404** can be configured to return a reference **7418** such as “atsign://claim2.alice@trustedservice” which can be provided to the requester. The requester can parse the reference **7418** and then perform a secondary lookup **7420** which can be provided to the trusted entity **7412**. The trusted entity **7412** can return a response **7422** to the requester. The response **7422** can be a value or can be another reference (e.g., to establish a chain of references) that can be further resolved. The response **7422** can be polymorphically determined, e.g., having a value determined based on an identity of the requester.

(591) In some implementations, a response to the lookup **7416** may include more than one third party reference that the requestor can choose from or match against. For example, the secondary server **7404** can return, in response to the lookup **7416**, a response of [“atsign://claim2.alice@service1”, “atsign://claim2.alice@service2”, “atsign://claim2.alice@service3”]. The requester can choose a given attester/service provider over other(s), should one service provider be preferable to (e.g., more trusted by) the requester.

(592) The @protocol can support anonymous attestations. For example, the @protocol can support providing proof of some claim (e.g., an attestation) while the claimant remains anonymous to the receiving entity/requester. For example, the trusted entity **7412** can provide, in some context, an anonymous response **7424** to an authorized requesting entity without divulging the @sign associated with the anonymous response **7424**. For example, within some context, an attestation of a “credit score”; e.g., 800 for some user has been validated by the trusted entity **7412** without exposing anything that could identify the user associated with the credit score. Such information might be very useful for a receiving entity to decide how best to respond or proceed in the given context. Anonymous attestations can be a portion of an Internet permissioning system that can simultaneously provide both accountability and anonymity on the Internet. Anonymous attestations can be implemented using two features of the @protocol: 1) the @protocol scheme (“atsign://”) which allows for the expression of a URL endpoint (e.g. “atsign://creditscore@trustedservice”) which might be an attestation; and 2) the polymorphic nature of @protocol responses.

(593) Attestations can be used for various other purposes. For example, attestation can be used for @protocol-enabled application certification. For example, an application developer can publish a checksum for an application/revision distribution package as a self-attestation that can be looked up and evaluated. For example, a lookup of e.g., myappversion5.8checksum@appdeveloper can return a checksum value, which can be compared to a checksum computed for received application code, to ensure that checksums match and that checksum or application code has not been altered.

(594) The preceding figures and accompanying description illustrate example processes and computer-implementable techniques. But system **100** (or its software or other components) contemplates using, implementing, or executing any suitable technique for performing these and other tasks. It will be understood that these processes are for illustration purposes only and that the described or similar techniques may be performed at any appropriate time, including concurrently,

individually, or in combination. In addition, many of the operations in these processes may take place simultaneously, concurrently, and/or in different orders than as shown. Moreover, system **100** may use processes with additional operations, fewer operations, and/or different operations, so long as the methods remain appropriate.

(595) In other words, although this disclosure has been described in terms of certain embodiments and generally associated methods, alterations and permutations of these embodiments and methods will be apparent to those skilled in the art. Accordingly, the above description of example embodiments does not define or constrain this disclosure. Other changes, substitutions, and alterations are also possible without departing from the spirit and scope of this disclosure.

Claims

1. A computer-implemented method comprising: receiving, at a server that manages access to a first decentralized resource directory of a providing entity in a system of decentralized resource directories, via a first connection and from a first requesting entity, an authentication request to authenticate the first requesting entity to the first decentralized resource directory, wherein the authentication request identifies the first requesting entity; generating, in response to the authentication request, an authentication challenge value and an authentication challenge key; providing, via the first connection and in response to the authentication request, an authentication challenge to the first requesting entity that instructs the first requesting entity to store the authentication challenge value for the authentication challenge key in a defined location in a second decentralized resource directory of the system of decentralized resource directories that is associated with the first requesting entity; receiving, via the first connection, a confirmation from the first requesting entity that the authentication challenge value has been stored by the first requesting entity for the authentication challenge key, in response to the authentication challenge, in the defined location in the second decentralized resource directory of the first requesting entity; establishing, in response to receiving the confirmation that the authentication challenge value has been stored for the authentication challenge key in the defined location in the second decentralized resource directory of the first requesting entity, a second connection with the second decentralized resource directory of the first requesting entity; sending, via the second connection, a first lookup request for a value stored for the authentication challenge key in the defined location in the second decentralized resource directory of the first requesting entity; receiving, via the second connection and from the second decentralized resource directory, a first response value in response to the first lookup request; comparing the first response value to the authentication challenge value to determine whether the first response value matches the authentication challenge value; and in response to determining that the first response value matches the authentication challenge value, responding to the authentication request, via the first connection, with an indication that the first requesting entity is authenticated to the first decentralized resource directory.

2. The computer-implemented method of claim 1, further comprising, in response to determining that the first response value does not match the authentication challenge value, dropping the first connection.

3. The computer-implemented method of claim 1, further comprising: receiving, via the first connection and from the first requesting entity, while the first requesting entity is authenticated to the first decentralized resource directory, a second lookup request for a value for a key stored in the first decentralized resource directory; determining that the first requesting entity has been provided access to a first stored value for the key; retrieving the first stored value for the key; and providing the first stored value as a second response value to the first requesting entity, in response to the second lookup request.

4. The computer-implemented method of claim 3, further comprising: receiving a third lookup request, via an unauthenticated third connection and from an unauthenticated second requesting

- entity, for a value for the key stored in the first decentralized resource directory; retrieving a publicly accessible value for the key; and providing the publicly accessible value for the key as a third response value, to the second requesting entity, in response to the third lookup request.
5. The computer-implemented method of claim 4, wherein the first stored value for the key provided to the authenticated first requesting entity is different from the publicly accessible value for the key.
6. The computer-implemented method of claim 4, further comprising: receiving, via a third connection and from a third requesting entity while the third requesting entity is authenticated to the first decentralized resource directory, a fourth lookup request for a value for the key stored in the first decentralized resource directory; determining that the third requesting entity has been provided access to a third stored value for the key, wherein the third stored value is different from the first stored value and the publicly accessible value for the key; retrieving the third stored value for the key; and providing the third stored value as a third response value to the third requesting entity, in response to the fourth lookup request.
7. The computer-implemented method of claim 3, further comprising: connecting, before sending the second lookup request, to a namespace directory to request connection information for the second decentralized resource directory; and using the connection information for the second decentralized resource directory to establish the second connection.
8. The computer-implemented method of claim 3, wherein the authentication request and the second lookup request are received from a user device associated with the first requesting entity.
9. The computer-implemented method of claim 3, wherein the authentication request and the second lookup request are received from a server that services requests for the second decentralized resource directory.
10. The computer-implemented method of claim 1, wherein the authentication challenge value is a randomly-generated identifier.
11. A non-transitory computer-readable storage medium coupled to one or more processors and having instructions stored thereon which, when executed by the one or more processors, cause the one or more processors to perform operations comprising: sending, via a first connection, an authentication request to authenticate as a requesting entity to a first decentralized resource directory of a providing entity, wherein the authentication request identifies the requesting entity and wherein the first decentralized resource directory is included in a system of decentralized resource directories; receiving, via the first connection, from the providing entity, and in response to the authentication request, an authentication challenge to store an authentication challenge value for an authentication challenge key in a defined location in a second decentralized resource directory of the system of decentralized resource directories that is associated with the requesting entity; storing, in response to the authentication challenge, in the defined location in the second decentralized resource directory, the authentication challenge value for the authentication challenge key; sending, in response to the authentication challenge, an authentication challenge response to the providing entity requesting the providing entity to verify the authentication challenge; and receiving an indication from the providing entity that the requesting entity is authenticated to the first decentralized resource directory as the requesting entity.
12. The computer-readable storage medium of claim 11, wherein the operations further comprise: sending, while authenticated to the first decentralized resource directory and via the first connection, a first lookup request to the first decentralized resource directory for a value for a key stored in the first decentralized resource directory; and receiving, in response to the first lookup request, a first value for the key.
13. The computer-readable storage medium of claim 12, wherein the operations further comprise: sending, as an unauthenticated requesting entity, a second lookup request, via a second connection, to the first decentralized resource directory for the value for the key stored in the first decentralized resource directory; and receiving, in response to the second lookup request, a second value for the

key, wherein the second value is different from the first value.

14. The computer-readable storage medium of claim 11, wherein the operations further comprise: connecting, before sending the authentication request, to a namespace directory to request connection information for the first decentralized resource directory; and using the connection information for the first decentralized resource directory to establish the first connection.

15. The computer-readable storage medium of claim 11, wherein the authentication challenge value is a randomly-generated identifier generated by the providing entity.
