

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12407590
Kind Code	B1
Date of Patent	September 02, 2025
Inventor(s)	Ainsworth; John et al.

Abandoned machines

Abstract

A system and method for obtaining, using one or more processors, network telemetry data; identifying, using the one or more processors, a set of active machines based on network telemetry data; analyzing, using the one or more processors, traffic associated with the set of active machines; determining, using the one or more processors, one or more abandoned machines in the set of active machines; and outputting, using the one or more processors, an identification of the one or more abandoned machines.

Inventors: Ainsworth; John (Burlington, MA), Minarik; Pavel (Burlington, MA), Knapek; Jiri (Burlington, MA)

Applicant: Progress Software Corporation (Burlington, MA)

Family ID: 96882264

Assignee: Progress Software Corporation (Burlington, MA)

Appl. No.: 18/522014

Filed: November 28, 2023

Publication Classification

Int. Cl.: H04L43/065 (20220101); H04L43/02 (20220101)

U.S. Cl.:

CPC H04L43/065 (20130101); H04L43/02 (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2015/0271192	12/2014	Crowley	726/22	H04L 41/149
2017/0147372	12/2016	Cudak	N/A	G06F 9/45558
2017/0359246	12/2016	Bodas	N/A	H04L 67/10
2020/0028871	12/2019	Thayer	N/A	H04L 63/1466
2020/0379529	12/2019	Le Goff	N/A	G06F 1/206
2022/0263725	12/2021	Mayer Zu Eissen	N/A	H04L 41/147
2023/0143232	12/2022	Balazs	709/223	H04L 63/1466
2023/0164043	12/2022	Sirov	709/224	H04L 41/5032

Primary Examiner: Joo; Joshua

Attorney, Agent or Firm: PATENT LAW WORKS LLP

Background/Summary

FIELD OF INVENTION

(1) The present disclosure relates to identification of one or more active but abandoned machines.

BACKGROUND

(2) It is not uncommon for a user to deploy a machine (virtual or physical), use it for a period, and then abandon (i.e., cease using) that machine without shutting or deactivating that machine.

SUMMARY

(3) In general, an innovative aspect of the subject matter described in this disclosure may be embodied in methods that include: obtaining, using one or more processors, network telemetry data; identifying, using the one or more processors, a set of active machines based on network telemetry data; analyzing, using the one or more processors, network telemetry data associated with the set of active machines; determining, using the one or more processors, one or more abandoned machines in the set of active machines; and outputting, using the one or more processors, an identification of the one or more abandoned machines.

(4) According to another innovative aspect of the subject matter described in this disclosure, a method to: obtain network telemetry data; identify a set of active machines based on network telemetry data; analyze the network telemetry data associated with the set of active machines; determining one or more abandoned machines in the set of active machines; and output an identification of the one or more abandoned machines.

(5) Other implementations of one or more of these aspects include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices. These and other implementations may each optionally include one or more of the following features.

(6) For instance, the features further include generating a set of machines to monitor for abandonment based on a set of machines known to a network traffic monitor, wherein the network telemetry data obtained is a subset of available network telemetry data available and is that network telemetry data that is associated with the set of machines to monitor for abandonment. For instance, the features further include determining a set of monitored machines, the set of monitored machines being monitored by a network traffic monitor that generates the network telemetry data,

and the set of monitored machines is further being monitored for abandonment, wherein the network telemetry data obtained describes the set of one or more monitored machines and includes the set of active machines. For instance, the features further include that the set of monitored machines is identified by an identifier in the network telemetry data, wherein the identifier includes one or more of a range of a MAC addresses, a range of IP addresses, a list of MAC addresses, a list of IP addresses, a universal unique identifier of a virtual machine, and an interface ID of network adapter. For instance, the features further include analyzing the network telemetry data associated with the set of active machines determining whether the network telemetry data includes one or more of a first characteristic associated with an active but unabandoned machine and a second characteristic associated with an abandoned machine. For instance, the features further include that the second characteristic associated with an abandoned machine is associated with one or more of a port number associated with traffic and a periodicity of the traffic, the traffic described by the network telemetry data. For instance, the features further include that determining the one or more abandoned machines in the set of active machines includes: analyzing the network telemetry data; determining that the network telemetry data associated with a first machine is uncharacteristic of human use; and identifying the first machine as one of the one or more abandoned machines. For instance, the features further include that the set of active machines includes one or more of a virtual machine and a physical machine. For instance, the features further include deactivating a first abandoned machine in the one or more identified, abandoned machines, wherein the deactivation is one of: automatic and based on the identification of the first abandoned machine as abandoned, or based on the identification of the first abandoned machine as abandoned and responsive to user input requesting abandonment. For instance, the features further include that the identification of the one or more abandoned machines is output in a report presented to a user.

(7) It should be understood that this list of features and advantages is not all-inclusive and many additional features and advantages are contemplated and fall within the scope of the present disclosure. Moreover, it should be understood that the language used in the present disclosure has been principally selected for readability and instructional purposes, and not to limit the scope of the subject matter disclosed herein.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The disclosure is illustrated by way of example, and not by way of limitation in the figures of the accompanying drawings in which like reference numerals are used to refer to similar elements.
- (2) FIG. 1 is a block diagram illustrating an example system for identifying one or more abandoned machines in accordance with some implementations.
- (3) FIG. 2 is a block diagram illustrating an example server in accordance with some implementations.
- (4) FIG. 3 is a block diagram of an abandoned machines identifier in accordance with some implementations.
- (5) FIG. 4 is a flowchart of an example method for identifying one or more abandoned machines in accordance with some implementations.

DETAILED DESCRIPTION

(6) It is not uncommon for a user to deploy a machine, use it for a period, and then abandon (i.e., cease using) that machine. Abandoned machines, while not being actively used by a user, are still active machines. Accordingly, those abandoned machines use resources and may present security exposures. For example, those abandoned machines may consume one or more of monetary resources (e.g., as an operational or capital expenditure), computation resources (e.g., bandwidth, processor cycles, memory, etc.), and physical resources (e.g., electricity, etc.). As another example,

the abandoned machines may negatively impact security exposure, e.g., by providing forgotten tunnels through firewalls or other security mechanisms that may be coopted by a nefarious user. Therefore, it may be beneficial to identify and deactivate abandoned machines. However, identifying abandoned machines is non-trivial. A user that initially deployed the machine and subsequently abandoned the machine may not realize that machine is still active and is likely unable to identify which machine he/she deployed and abandoned even if presented a list of active machines. For example, machines may be spooled up by developers in a development environment and be left abandoned, and the developer may not be aware of which machines he/she spooled up or whether they were deactivated or remain active and abandoned. The systems and methods described herein may beneficially identify abandoned machines, which may, depending on the implementation, be manually and/or automatic deactivated.

(7) FIG. 1 is a block diagram illustrating an example system **100** for identifying one or more abandoned machines according to one implementation. The illustrated system **100** includes a client device **106**, one or more machines **132a . . . n**, and a server **122**, which are communicatively coupled via a network **102** for interaction with one another. For example, the client devices **106** may be coupled to the network **102** via signal line **104** and may be accessed by user **112** as illustrated by line **110**. The machines **132a . . . n** may be coupled to the network **102** via signal lines **118a . . . n**. The server **122** may be coupled to the network **102** via signal line **116**. The use of the nomenclature “a” and “n” in the reference numbers indicates that any number of those elements having that nomenclature may be included in the system **100**.

(8) The network **102** may include any number of networks and/or network types. For example, the network **102** may include, but is not limited to, one or more local area networks (LANs), wide area networks (WANs) (e.g., the Internet), virtual private networks (VPNs), mobile networks (e.g., the cellular network), wireless wide area network (WWANs), Wi-Fi networks, WiMAX® networks, Bluetooth® communication networks, peer-to-peer networks, other interconnected data paths across which multiple devices may communicate, various combinations thereof, etc. Data transmitted by the network **102** may include packetized data (e.g., Internet Protocol (IP) data packets) that is routed to designated computing devices coupled to the network **102**. In some implementations, the network **102** may include a combination of wired and wireless (e.g., terrestrial or satellite-based transceivers) networking software and/or hardware that interconnects the computing devices of the system **100**. For example, the network **102** may include packet-switching devices that route the data packets to the various computing devices based on information included in a header of the data packets.

(9) The data exchanged over the network **102** can be represented using technologies and/or formats including the hypertext markup language (HTML), the extensible markup language (XML), JavaScript Object Notation (JSON), Binary JavaScript Object Notation (BSON), Comma Separated Values (CSV), etc. In addition, all or some of links can be encrypted using conventional encryption technologies, for example, the secure sockets layer (SSL), Secure HTTP (HTTPS) and/or virtual private networks (VPNs) or Internet Protocol security (IPsec). In another implementation, the entities can use custom and/or dedicated data communications technologies instead of, or in addition to, the ones described above. Depending upon the implementation, the network **102** can also include links to other networks.

(10) The client devices **106** are computing devices having data processing and communication capabilities. While FIG. 1 illustrates one client devices **106**, the present specification may apply to any system architecture having one or more client devices **106**. In some implementations, a client device **106** may include a processor (e.g., virtual, physical, etc.), a memory, a power source, a network interface, and may include other components whether software or hardware, such as a display, graphics processor, wireless transceivers, keyboard, camera, sensors, firmware, operating systems, drivers, various physical connection interfaces (e.g., USB, HDMI, etc.). The client devices **106a . . . 106n** may couple to and communicate with one another and the other entities of the

system **100** via the network **102** using a wireless and/or wired connection.

(11) Examples of client devices **106** may include, but are not limited to, mobile phones (e.g., feature phones, smart phones, etc.), tablets, laptops, desktops, netbooks, server appliances, servers, virtual machines, TVs, set-top boxes, media streaming devices, portable media players, navigation devices, personal digital assistants, etc. While one client device **106** is depicted in FIG. **1**, the system **100** may include any number of client devices **106**. For example, a first user **112** may use a first client device **106** to use a machine **132a** and a second user **112b** may use a second (not shown) client device **106** to access an identification of the one or more abandoned machines generated by the abandoned machine identifier **220**. In addition, the different instances of the client devices **106** may be the same or different types of computing devices.

(12) Depending on the implementation and use case, a client device **106** may present a set of abandoned machines identified by the abandoned machines IDer **220**, e.g., in a report. In some implementations, the client device **106** may be used by a developer, e.g., to request and spool up one or more of the machines **132a . . . n** which may be potentially become abandoned.

(13) The server **122** may include one or more computing devices having data processing, storing, and communication capabilities. For example, the server **122** may include one or more hardware servers, server arrays, storage devices, systems, etc., and/or may be centralized or distributed/cloud-based. In some implementations, the server **122** may include one or more virtual servers, which operate in a host server environment and access the physical hardware of the host server including, for example, a processor, memory, storage, network interfaces, etc., via an abstraction layer (e.g., a virtual machine manager).

(14) In some implementations, the system **100** includes a plurality of network telemetry traffic monitors (not shown) that generate network telemetry data. For example, while not shown, the network **102** may include network traffic telemetry monitors (not shown) including any of the following (individually or in combination): one or more probes, one or more routers, one or more switches, one or more network device that generate or provide telemetry data, one or more cloud platforms that generate or provide telemetry data, etc.

(15) In some implementations, network telemetry data is collected or aggregated and used by the abandoned machine identifier **220** to identify one or more abandoned machines. In some implementations, the collector, or aggregator, of network telemetry data is referred to herein as the network traffic monitor **142**. Examples of a network traffic monitor **142** may include a Progress Software Corporation's Flowmon® collector; however, other examples exist and are within the scope of this disclosure. In one implementation, the server **122** includes an instance of a network traffic monitor **142** and an abandoned machines identifier **220** that uses the network telemetry data generated by the network traffic monitor **142**.

(16) The one or more machines **132a . . . 132n** may include, or represent, one or more computing devices having data processing, storing, and communication capabilities. For example, the machine(s) **132** may include, or represent, one or more of a physical/hardware server(s), server array(s), storage device(s), systems, appliance(s), etc., or any combination thereof, and/or may be centralized or distributed/cloud-based. In some implementations, the one or more machines **132** may include one or more virtual machines. For example one or more servers, which operate in a host server environment and access the physical hardware of the host server including, for example, a processor, memory, storage, network interfaces, etc., via an abstraction layer (e.g., a virtual machine manager). In some implementations, the one or more machines **132** may include a mixture of machine types including at least one hardware and at least one virtual machines.

(17) It should be understood that the system **100** illustrated in FIG. **1** is representative of an example system for identifying one or more abandoned machines according to one implementation and that a variety of different system environments and configurations are contemplated and are within the scope of the present disclosure. For instance, various functionality may be moved from a server (as illustrated) to a machine or a client, or vice versa and some implementations may include

additional or fewer computing devices, services, and/or networks, and may implement various functionality client or server-side. Further, various entities of the system **100** may be integrated into to a single computing device or system or additional computing devices or systems, etc.

(18) FIG. 2 is a block diagram of an example server **122** according to one implementation. The server **122**, as illustrated, may include a processor **202**, a memory **204** and a communication unit **208**, which may be communicatively coupled by a communications bus **206**. The server **122** depicted in FIG. 2 is provided by way of example and it should be understood that it may take other forms and include additional or fewer components without departing from the scope of the present disclosure. For example, the server **122** may include a storage device, input devices **212** and/or output devices **214** (e.g., a display, a keyboard, a mouse, touch screen, speakers, etc.), various operating systems, sensors, additional processors, and other physical configurations. Additionally, it should be understood that the computer architecture depicted in FIG. 2 and described herein with reference to the server **122** may be applied to other entities in a system **100**, for example, a client device **106** or a machine **132**.

(19) The processor **202** may execute code, routines and software instructions by performing various input/output, logical, and/or mathematical operations. The processor **202** have various computing architectures to process data signals including, for example, a complex instruction set computer (CISC) architecture, a reduced instruction set computer (RISC) architecture, and/or an architecture implementing a combination of instruction sets. The processor **202** may be physical and/or virtual and may include a single core or plurality of processing units and/or cores. In some implementations, the processor **202** may be capable of generating and providing electronic display signals to a display device (not shown), supporting the display of images, capturing and transmitting images, performing complex tasks including various types of feature extraction and sampling, etc. In some implementations, the processor **202** may be coupled to the memory **204** via the bus **206** to access data and instructions therefrom and store data therein. The bus **206** may couple the processor **202** to the other components of the server **122** including, for example, the memory **204** and communication unit **208**.

(20) The memory **204** may store and provide access to data to the other components of the server **122**. In some implementations, the memory **204** may store instructions and/or data that may be executed by the processor **202**. For example, in the illustrated implementation, the memory **204** may store one or more of the abandoned machines identifier **220** and the network traffic monitor **142**. The memory **204** is also capable of storing other instructions and data, including, for example, an operating system, hardware drivers, other software applications, databases, etc. The memory **204** may be coupled to the bus **206** for communication with the processor **202** and the other components of the server **122**. In the illustrated implementation, the memory **204** stores an abandoned machines identifier **220** and a network traffic monitor **142**. In one implementation, the network traffic monitor **142** generates telemetry data, which is used by abandoned machines identifier **220** to one or more abandoned machines.

(21) The memory **204** includes a non-transitory computer-usable (e.g., readable, writeable, etc.) medium, which can be any apparatus or device that can contain, store, communicate, propagate or transport instructions, data, computer programs, software, code, routines, etc., for processing by or in connection with the processor **202**. In some implementations, the memory **204** may include one or more of volatile memory and non-volatile memory. For example, the memory **204** may include, but is not limited, to one or more of a dynamic random-access memory (DRAM) device, a static random access memory (SRAM) device, a discrete memory device (e.g., a PROM, EPROM, ROM), a hard disk drive, an optical disk drive (CD, DVD, Blue-ray™, etc.). It should be understood that the memory **204** may be a single device or may include multiple types of devices and configurations.

(22) The bus **206** can include a communication bus for transferring data between components of a server **122** and/or between computing devices (e.g. between one or more of the client device **106**

and non-relational database **120**), a network bus system including the network **102** or portions thereof, a processor mesh, a combination thereof, etc. In some implementations, the schema tool **109**, its sub-components and various other software operating on the computing device **106/122** (e.g., an operating system, etc.) may cooperate and communicate via a software communication mechanism implemented in association with the bus **206**. The software communication mechanism can include and/or facilitate, for example, inter-process communication, local function or procedure calls, remote procedure calls, an object broker (e.g., CORBA), direct socket communication (e.g., TCP/IP sockets) among software modules, UDP broadcasts and receipts, HTTP connections, etc. Further, any or all of the communication could be secure (e.g., SSH, HTTPS, etc.).

(23) The communication unit **208** may include one or more interface devices (I/F) for wired and/or wireless connectivity with the network **102**. For instance, the communication unit **208** may include, but is not limited to, CAT-type interfaces; wireless transceivers for sending and receiving signals using radio transceivers (4G, 3G, 2G, etc.) for communication with the mobile network **102**, and radio transceivers for Wi-Fi™ and close-proximity (e.g., Bluetooth®, NFC, etc.) connectivity, etc.; USB interfaces; various combinations thereof; etc. In some implementations, the communication unit **208** can link the processor **202** to the network **102**, which may in turn be coupled to other processing systems. The communication unit **208** can provide other connections to the network **102** and to other entities of the system **100** using various standard network communication protocols.

(24) As mentioned above, the server **122** may include other and/or fewer components. Examples of other components may optionally include one or more of an output device **214**, an input device **212**, a sensor (not shown), etc. In one implementation, the server **122** includes an output device **214** including a display. The display may display electronic images and data for presentation to a user **112**. The display may include any display device, monitor or screen, including, for example, an organic light-emitting diode (OLED) display, a liquid crystal display (LCD), etc. In some implementations, the display may be a touch-screen display capable of receiving input from a stylus, one or more fingers of a user **112**, etc. For example, the display may be a capacitive touch-screen display capable of detecting and interpreting multiple points of contact with the display surface.

(25) The optional input device **212** may include any device for inputting information into the server **122**. In some implementations, the input device **212** may include one or more peripheral devices. For example, the input device may include a keyboard (e.g., a QWERTY keyboard or keyboard in any other language), a pointing device (e.g., a mouse or touchpad), microphone, an image/video capture device (e.g., camera), etc. In some implementations, the input device may include a touch-screen display capable of receiving input from the one or more fingers of the user **112**. For example, the user **112** could interact with an emulated (i.e., virtual or soft) keyboard displayed on the touch-screen display by using fingers to contacting the display in the keyboard regions.

(26) Example Abandoned Machines Identifier **220**

(27) Referring now to FIG. **3**, the abandoned machines identifier **220** is shown in more detail according to one implementation. The abandoned machines identifier **220** identifies one or more abandoned machines. In one implementation, the abandoned machines identifier **220** comprises an abandoned machine profile editor **302**, a network telemetry data obtainer **304**, an active machine identifier **306**, a telemetry analyzer **308**, a reporting engine, and an abandoned machine deactivator **312**. It will be recognized that the components **302**, **304**, **306**, **308**, **310**, and **312** comprised in the abandoned machines identifier **220** are not necessarily all on the same server **122**. In one implementation, the components **302**, **304**, **306**, **308**, **310**, and **312** and/or their functionality are distributed across multiple servers **122**.

(28) The abandoned machine profile editor **302** includes code and routines for generating an abandoned machine profile. In one implementation, the abandoned machine profile editor **302** is a set of instructions executable by the processor **202**. In another implementation, the abandoned machine profile editor **302** is stored in the memory **204** and is accessible and executable by the

processor **202**. In either implementation, the abandoned machine profile editor **302** is adapted for cooperation and communication with the processor **202**, other components of the server **122** and other components of the abandoned machines identifier **220**.

(29) In some implementations, the abandoned machine profile editor **302** generates an abandoned machine profile during a configuration period and the abandoned machine profile is used during runtime by one or more other components **304, 306, 308, 310, 312** of the abandoned machines identifier **220**. The configurability of the abandoned machine profile may allow customization to different machine environments. For example, the abandoned machine profile editor **302** may generate a different profile for a developer environment (e.g., to identify abandoned machines used by a team of developers developing an application) than for a production environment (e.g., to identify abandoned machines used by customers/end-users).

(30) In some implementations, an abandoned machine profile includes an identification of a set of monitored machines and a traffic definition. The set of monitored machines identifies one or more machines monitored, by the abandoned machines identifier **220**, for abandonment. In some implementations, the set of monitored machines is identified based on identifiers present in the network telemetry data. Depending on the implementation and use case, the identification of the set of monitored, by the abandoned machines identifier **220**, machines is based on one or more of an IP address, MAC address, a universally unique identifier (UUID) such as that of a virtual machine, an interface ID of a network adapter, or other identifier. For example, the identification of the set of monitored machines may comprise one or more of a range of MAC addresses, a range of IP addresses, a list of MAC addresses, a list of IP addresses, a range or list of UUIDs of virtual machines, a range or list of interface ID of a network adapters, etc. As another example, the set of monitored machines may include and be identified by an IP subnet such as 192.168.50.0/23.

(31) In some implementations, the identified set of one or more machines, which are monitored by the abandoned machines identifier **220** for abandonment, are also monitored by the network traffic monitor **142**. In some implementations, the abandoned machine profile editor **302** may be communicatively coupled to the network traffic monitor **142** to generate a set of suggestions or provide a list of known (to the network traffic monitor **142**) machines which may be selected to generate the set of monitored machines.

(32) In some implementations, the set of one or more machines monitored for abandonment, as defined in a particular profile, may be a subset of those machines available to be monitored for abandonment, e.g., the one or more machines may be specific to only those machines used by a set of developers associated with an entity. It should be understood that by generating multiple profiles using the abandoned machine profile editor **302**, the systems and methods described herein may provide more accurate identification of abandoned machines (e.g., by better tailoring definitions of abandonment to different sets of machines) and more granular control (e.g., over whether to err on the side of over inclusivity or under inclusivity, or whether to auto-deactivate identified abandoned machines, etc.). For example, the profiles may be defined such that machines associated with the production environment are not automatically deactivated and/or are subject to a more stringent definition of what it means to be abandoned to minimize the risk of interruptions to users (e.g., customers) of the machine, should that machine be inaccurately identified as abandoned.

(33) The traffic definition defines one or more characteristics describing traffic of an active machine and/or an abandoned machine. For example, the traffic definition may include a set of rules including presence of traffic on a particular port number (i.e., a first characteristic that may indicate active status), periodicity of the traffic (i.e., a second characteristic that, if regular, could represent regular system updates rather than active use by a user), a quantity of the traffic (i.e. a third characteristic that if relatively small may be indicative of disuse, particular if periodic, as may be the case with a heartbeat signal or some sort of system polling), a source or destination of the traffic (i.e., a fourth or fifth characteristic that may be indicative of abandonment/disuse, e.g., traffic between the machine owner/operator, such as AWS or its load balancer, vs to a user's client

device), etc.

(34) The one or more characteristics in the traffic definition may vary. In some implementations, the one or more characteristics are a subset of the attributes available in network telemetry data obtainable from the network traffic monitor **142**. The variety of characteristics and network telemetry data attributes are so varied and numerous as to provide a complete listing. However, it should be recognized that the number and variety of available characteristics provide the systems and methods herein with a great degree of flexibility and granularity to tailor and optimize the traffic definition. For example, to reduce false positives, to reduce false negatives, to reduce the number of characteristics (or network telemetry data attributes) obtained and used to evaluate for abandonment, etc.

(35) In some implementations, the profile may define one or more times. For example, the profile may identify a frequency at which the profile is to be applied by the abandoned machines identifier **220** and submodules thereof for abandonment (e.g., every X seconds, minutes, hours, days, weeks, months, years, etc.) to identify whether any machines are abandoned. As another example, the definition may define a time related to the network traffic definition, such as machine is abandoned if the machine has not received management traffic on port **443** or **22** for X days.

(36) Depending on the implementation, the traffic definition may one or more of positively identify an abandoned machine (i.e., the definition describes one or more characteristics that, when met, indicate abandonment of the associated machine), negatively identify an abandoned machine (i.e., the definition describes one or more characteristics of an active and in-use/unabandoned machine, and an abandoned machine fails to satisfy the one or more characteristics), or a combination thereof (e.g., use a first portion of a definition to identify unabandoned, active machines and then apply a second portion of a definition to identify which of the remaining machines are abandoned) based on their associated traffic.

(37) Depending on the implementation, the abandonment machine profile editor **302** may generate an abandoned machine profile using one or more of a variety of mechanisms. In some implementations, an abandoned machine profile may be generated through user input, e.g., via an input device such as a keyboard and/or pointer. In some implementations, the abandoned machine profile editor **302** may generate an abandoned machine profile using one or more of an API, script, and UI. For example, in some implementations, abandonment machine profile editor **302** may generate an abandoned machine profile by presenting a UI during configuration and receiving user input defining the set of one or more machines for abandonment (e.g., via drag-and-drop from a list of machines monitored by the network traffic monitor **142**) and a traffic definition.

(38) In some implementations, the abandoned machine profile editor **302** may at least partially generate an abandoned machine profile using machine learning. In some implementations, the abandoned machine profile editor **302** may train, validate, and test a machine learning model to classify a set of machines associated with an entity so that each class of machine may be associated with its own abandoned machine profile. For example, the abandoned machine profile editor **302** may automatically generate a profile associated with a set of machines used exclusively by a team of developers and populate the machine identifiers associated with that set of one or more machines to monitor for abandonment in association with that profile. As another example, the abandoned machine profile editor **302** may identify a set of machines that are frequently, or have historically, been abandoned and automatically generate a profile associated with that set of machines. In some implementations, the abandoned machine profile editor **302** may train, validate, and test a machine learning model to determine one or more of the characteristics of an abandoned and/or of an unabandoned machine and populate the traffic definition based on that model. For example, in some implementations, the traffic definition may include the model as the whole or partial definition. As another example, in some implementations, the model may modify the characteristics or sub-characteristics used in the traffic definition (e.g., modify a weight assigned to each characteristic, modify a threshold of what a “small” amount of data means, add or remove a port

from a port list, etc.).

(39) The abandoned machine profile editor **302** may train one or more abandoned machine identification models using a variety of machine learning techniques, depending on the implementation and use case, including supervised learning, unsupervised learning, semi-supervised learning, etc. The varieties of supervised, semi-supervised, and unsupervised machine learning algorithms that may be used, by the abandoned machine profile editor **302**, to train the one or more abandoned machine identification models are so numerous as to defy a complete list. Example algorithms include, but are not limited to, a decision tree; a gradient boosted tree; boosted stumps; a random forest; a support vector machine; a neural network; a recurrent neural network; a recurrent neural network; deep learning; long short-term memory; transformer; logistic regression (with regularization), linear regression (with regularization); stacking; a Markov model; Markov chain; support vector machines; and others. In some implementations, the abandoned machine identification model may include a binary classifier to label a machine as abandoned or not abandoned.

(40) The training, training data, validation, and validation data may vary depending on the implementation. In some implementations, the abandoned machines profile editor **302** may be retrained based on user input. For example, when a user selects to manually deactivate a machine identified by the system and methods herein as abandoned, that may be used to retrain, or through reinforced learning, to improve the accuracy of the model.

(41) In some implementations, abandoned machine profile editor **302** passes the abandoned machine profile(s) to one or more of the network telemetry data obtainer **304**, the active machine identifier **306**, and the telemetry analyzer **308**. In some implementations, the abandoned machine profile editor **302** stores the abandoned machine profile(s) in memory **204** (or any other non-transitory storage medium communicatively accessible), and the other modules of the abandoned machines identifier **220** (e.g., one or more of the network telemetry data obtainer **304**, the active machine identifier **306**, and the telemetry analyzer **308**) may retrieve the abandoned machine profile(s) from the memory **204** (or other non-transitory storage medium).

(42) The network telemetry data obtainer **304** includes code and routines for obtaining network telemetry data. In one implementation, the network telemetry data obtainer **304** is a set of instructions executable by the processor **202**. In another implementation, the network telemetry data obtainer **304** is stored in the memory **204** and is accessible and executable by the processor **202**. In either implementation, the network telemetry data obtainer **304** is adapted for cooperation and communication with the processor **202**, other components of the server **122** and other components of the abandoned machine identifier **220**.

(43) The network telemetry data obtainer **304** obtains network telemetry data. In one implementation, the network telemetry data obtainer **304** obtains network telemetry data from, or stored by, the network traffic monitor **142**. In some implementations, the network traffic monitor **142** may include Progress Software Corporation's Flowmon®; however, Flowmon® is merely one example of a network traffic monitor **142** and others exist and may be used. For example, in some implementations, the abandoned machines IDer **220** is network traffic monitor **142** agnostic.

(44) Depending on the implementation and use case the network telemetry data obtained by the network telemetry data obtainer **304** may vary in the content of what the obtained network telemetry data represents and/or the format in which the network telemetry data is stored. In some implementations, the network telemetry data obtainer **304** may obtain only a subset of network telemetry data available to the network telemetry data obtainer **304**. For example, in some implementations, the network telemetry data obtainer **304** may obtain a subset of network telemetry data based on one or more profiles (e.g., only network telemetry data describing network traffic associated with a particular set of machines associated with that profile and/or relating to the traffic definition, such as relating to a particular port). In some cases, different network traffic monitors **142** may monitor different aspects of network traffic and/or the same aspects differently, and, in

some implementations, the network telemetry data obtainer **304** is network traffic monitor **142** agnostic and/or configurable to obtain the relevant network telemetry data (i.e., the network telemetry data used by the other components of the abandoned machines identifier **220**, as described below) from different network traffic monitors **142**. In some implementations, the network traffic monitor **142** represents Progress Software Corporation's Flowmon®. Depending on the implementation and use case, the network telemetry data may vary in format. Examples of network telemetry data formats may include, but are not limited to, one or more of netflow, ipfix, jflow, vpc flow logs, sflow, etc. In some implementations, the network telemetry data obtainer **304** may obtain the telemetry data and determine a common representation (e.g., convert and store the telemetry data in a single, common format).

(45) In one implementation, the network telemetry data obtainer **304** passes the telemetry data to one or more modules of the abandoned machines identifier **220** (e.g., the active machine identifier **306** and/or the telemetry analyzer **308**). In one implementation, the network telemetry data obtainer **304** stores the telemetry data in memory **204** (or any other non-transitory storage medium communicatively accessible). The other modules of the abandoned machines identifier **220** may retrieve the telemetry data from the memory **204** (or other non-transitory storage medium).

(46) The active machine identifier **306** determines a set of active machines. In some implementations, the active machine identifier **306** identifies the set of machines to be monitored (e.g., an IP subnet from an abandoned machine profile), determines which of those machines are active (e.g., based on telemetry data accessible via the network telemetry obtainer **304**), and generates an active machine list. Depending on the implementation and use case, the set of active machines may include active-but-abandoned machines, active-but-unabandoned machines, or a combination thereof.

(47) In one implementation, the active machine identifier **306** passes the set of active machines to the telemetry analyzer **308**. In one implementation, the active machine identifier **306** stores the set of active machines in memory **204** (or any other non-transitory storage medium communicatively accessible). The other modules of the abandoned machine identifier **220** including, e.g., the telemetry analyzer **308**, may retrieve the set of active machines by accessing the memory **204** (or other non-transitory storage medium).

(48) The telemetry analyzer **308** includes code and routines for determining which active machines are abandoned or unabandoned depending on the implementation. In one implementation, the telemetry analyzer **308** is a set of instructions executable by the processor **202**. In another implementation, the telemetry analyzer **308** is stored in the memory **204** and is accessible and executable by the processor **202**. In either implementation, the telemetry analyzer **308** is adapted for cooperation and communication with the processor **202**, other components of the computing device **106/122** (e.g., the network traffic monitor **142**) and other components of the abandoned machines identifier **220**.

(49) In some implementations, the telemetry analyzer **308** identifies which machines are abandoned based on whether the network telemetry data describing network traffic associated with a machine is characteristic or uncharacteristic of human machine use. In some implementations, the telemetry analyzer **308** applies a traffic definition, e.g., from an abandoned machines profile, to determine whether network telemetry data associated with an active machine is characteristic of human use, automated use, or both, and thereby determine whether a monitored and active machine is abandoned or in-use.

(50) In some implementations, the telemetry analyzer **308** passes identification of the abandoned machine(s) to one or more of the reporting engine **310** and the abandoned machines deactivator **312**. In one implementation, the telemetry analyzer **308** stores the identification of the abandoned machine(s) in memory **204** (or any other non-transitory storage medium communicatively accessible). The other modules of the abandoned machines identifier **220** including, e.g., the reporting engine **310** and/or the abandoned machines deactivator **312**, may retrieve the

identification of the abandoned machine(s) by accessing the memory **204** (or other non-transitory storage medium).

(51) The reporting engine **310** includes code and routines for generating and presenting a report. In one implementation, the reporting engine **310** is a set of instructions executable by the processor **202**. In another implementation, the reporting engine **310** is stored in the memory **204** and is accessible and executable by the processor **202**. In either implementation, the reporting engine **310** is adapted for cooperation and communication with the processor **202**, other components of the server **122** and other components of the abandoned machines identifier **220**.

(52) In one implementation, reporting engine **310** generates and presents a report. In some implementations, the report identifies the one or more active, but abandoned machines. For example, the report provides a list of machine identifiers, such as IP and/or MAC addresses, associated with each abandoned machine. In some implementations, the presentation may include presentation to a human (e.g., a system administrator) for review. In some implementations, the presentation may include presentation to the abandoned machine deactivator **312** for deactivation of those machines identified as abandoned, whether automatically or subject to human review and approval.

(53) In some implementations, the reporting engine **310** may obtain, provide, and track supplemental information. For example, when machine **132a** is determined to be abandoned, the reporting engine **310** may obtain and provide supplemental information describing one or more of when machine **132a** was deployed, what user initiated the machine's deployment, when the machine **132a** was last used, which user was the last to use machine **132a**, a user group associated with the last user, a user group associated with the user that initiated the machine's deployment, etc. Such information may be useful to identify users or user groups that may benefit from education or intervention due to patterns of abandonment (e.g., frequent abandonment). It should be recognized that the preceding are merely examples of supplemental information and others exist, are contemplated, and may be used without departing from this disclosure.

(54) In some implementations, the abandoned machines identifier **220** may omit an abandoned machine deactivator **312**. For example, once the abandoned machines identifier **220** identifies and presents the one or more abandoned machines, it is up to a user **112** to act (e.g., to deactivate a machine identified as abandoned). In some implementations, the deactivation actions and/or mechanism may be outside the scope of the abandoned machines identifier **220**, e.g., the user physically manipulates a power switch on an abandoned machine, or using a separate mechanism, application, or interface to deactivate the abandoned machine. In some implementations, deactivations or a subset of deactivations may be performed by the abandoned machine deactivator **312** of an abandoned machines identifier **220**.

(55) In some implementations, the reporting engine **310** passes the report to the abandoned machine deactivator **312**. For example, the reporting engine **310** is communicatively coupled to the abandoned machine deactivator **312** to send the report. In some implementations, the reporting engine **310** stores the report in memory **204** (or any other non-transitory storage medium communicatively accessible), and the abandoned machine deactivator **312** may retrieve the report by accessing the memory **204** (or other non-transitory storage medium).

(56) The abandoned machine deactivator **312** includes code and routines for deactivating one or more machines. In some implementations, the abandoned machine deactivator **312** is a set of instructions executable by the processor **202**. In another implementation, the abandoned machine deactivator **312** is stored in the memory **204** and is accessible and executable by the processor **202**. In either implementation, the abandoned machine deactivator **312** is adapted for cooperation and communication with the processor **202**, other components of the computing device **106/122** and other components of the abandoned machines identifier **220**.

(57) Depending on the implementation and use case, the abandoned machine deactivator **312** may deactivate abandoned machines with different levels of autonomy. For example, in some

implementations, the deactivation by the abandoned machine deactivator **312** is manual. For example, the abandoned machine deactivator **312** requires receipt of user input (e.g., user input explicitly instructing the abandoned machine deactivator **312** to deactivate a machine identified as abandoned) before deactivating a machine identified as abandoned and does not automatically deactivate an identified, abandoned machine without that user input/authorization. In some implementations, the deactivation performed by the abandoned machine deactivator **312** is automatic. For example, the abandoned machine deactivator **312** automatically deactivates a machine identified as abandoned responsive to identification of the machine as abandoned by the telemetry analyzer **308**.

(58) In some implementations, one or more of whether a report is generated, the format and/or content of a report, whether the abandoned machine(s) is deactivated manually or automatically may be specified in the abandoned machines profile via the abandoned machine profile editor. This may allow granular control of the definition of “abandoned” and treatment of machines identified as “abandoned” for different sets of machines. For example, a more inclusive definition of abandonment and auto-deactivation may be desired for internal environments to reduce costs (e.g., to rapidly deactivate machines abandoned by a company's development team), while a less inclusive definition and manual deactivation may be desired for machines that host applications used by customers to reduce the risk of a negative impacts on those customers.

(59) Example Methods

(60) FIG. 4 depicts method **400** performed by the system described above in reference to FIGS. 1-3. The method **400** begins at block **402**. At block **402**, the network telemetry data obtainer **304** obtains network telemetry data. At block **404**, the active machine identifier **306** identifies one or more active machines based on network telemetry data. At block **406**, the telemetry analyzer **308** analyzes the network telemetry data associated with active machines. At block **408**, the telemetry analyzer **308** determines which active machines are abandoned. At block **410**, the reporting engine **310** outputs an identification of the one or more abandoned machines.

(61) Other Considerations

(62) In the above description, for purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present disclosure. However, it should be understood that the technology described herein can be practiced without these specific details. Further, various systems, devices, and structures are shown in block diagram form in order to avoid obscuring the description. For instance, various implementations are described as having particular hardware, software, and user interfaces. However, the present disclosure applies to any type of computing device that can receive data and commands, and to any peripheral devices providing services.

(63) Reference in the specification to “one implementation” or “an implementation” means that a particular feature, structure, or characteristic described in connection with the implementation is included in at least one implementation. The appearances of the phrase “in one implementation” in various places in the specification are not necessarily all referring to the same implementation.

(64) In some instances, various implementations may be presented herein in terms of algorithms and symbolic representations of operations on data bits within a computer memory. An algorithm is here, and generally, conceived to be a self-consistent set of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, transferred, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

(65) It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. Unless specifically stated otherwise as apparent from the following discussion, it is

appreciated that throughout this disclosure, discussions utilizing terms including “processing,” “computing,” “calculating,” “determining,” “displaying,” or the like, refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage, transmission or display devices.

(66) Various implementations described herein may relate to an apparatus for performing the operations herein. This apparatus may be specially constructed for the required purposes, or it may comprise a general-purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program may be stored in a computer readable storage medium, including, but is not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, flash memories including USB keys with non-volatile memory or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

(67) The technology described herein can take the form of an entirely hardware implementation, an entirely software implementation, or implementations containing both hardware and software elements. For instance, the technology may be implemented in software, which includes but is not limited to firmware, resident software, microcode, etc.

(68) Furthermore, the technology can take the form of a computer program product accessible from a computer-usable or computer-readable medium providing program code for use by or in connection with a computer or any instruction execution system. For the purposes of this description, a computer-usable or computer readable medium can be any non-transitory storage apparatus that can contain, store, communicate, propagate, or transport the program for use by or in connection with the instruction execution system, apparatus, or device.

(69) A data processing system suitable for storing and/or executing program code may include at least one processor coupled directly or indirectly to memory elements through a system bus. The memory elements can include local memory employed during actual execution of the program code, bulk storage, and cache memories that provide temporary storage of at least some program code in order to reduce the number of times code must be retrieved from bulk storage during execution. Input/output or I/O devices (including but not limited to keyboards, displays, pointing devices, etc.) can be coupled to the system either directly or through intervening I/O controllers.

(70) Network adapters may also be coupled to the system to enable the data processing system to become coupled to other data processing systems, storage devices, remote printers, etc., through intervening private and/or public networks. Wireless (e.g., Wi-Fi™) transceivers, Ethernet adapters, and modems, are just a few examples of network adapters. The private and public networks may have any number of configurations and/or topologies. Data may be transmitted between these devices via the networks using a variety of different communication protocols including, for example, various Internet layer, transport layer, or application layer protocols. For example, data may be transmitted via the networks using transmission control protocol/Internet protocol (TCP/IP), user datagram protocol (UDP), transmission control protocol (TCP), hypertext transfer protocol (HTTP), secure hypertext transfer protocol (HTTPS), dynamic adaptive streaming over HTTP (DASH), real-time streaming protocol (RTSP), real-time transport protocol (RTP) and the real-time transport control protocol (RTCP), voice over Internet protocol (VOIP), file transfer protocol (FTP), WebSocket (WS), wireless access protocol (WAP), various messaging protocols (SMS, MMS, XMS, IMAP, SMTP, POP, WebDAV, etc.), or other known protocols.

(71) Finally, the structure, algorithms, and/or interfaces presented herein are not inherently related to any particular computer or other apparatus. Various general-purpose systems may be used with programs in accordance with the teachings herein, or it may prove convenient to construct more

specialized apparatus to perform the required method blocks. The required structure for a variety of these systems will appear from the description above. In addition, the specification is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages may be used to implement the teachings of the specification as described herein.

(72) The foregoing description has been presented for the purposes of illustration and description. It is not intended to be exhaustive or to limit the specification to the precise form disclosed. Many modifications and variations are possible in light of the above teaching. It is intended that the scope of the disclosure be limited not by this detailed description, but rather by the claims of this application. As will be understood by those familiar with the art, the specification may be embodied in other specific forms without departing from the spirit or essential characteristics thereof. Likewise, the particular naming and division of the modules, routines, features, attributes, methodologies and other aspects are not mandatory or significant, and the mechanisms that implement the specification or its features may have different names, divisions and/or formats.

(73) Furthermore, the modules, routines, features, attributes, methodologies and other aspects of the disclosure can be implemented as software, hardware, firmware, or any combination of the foregoing. Also, wherever a component, an example of which is a module, of the specification is implemented as software, the component can be implemented as a standalone program, as part of a larger program, as a plurality of separate programs, as a statically or dynamically linked library, as a kernel loadable module, as a device driver, and/or in every and any other way known now or in the future. Additionally, the disclosure is in no way limited to implementation in any specific programming language, or for any specific operating system or environment. Accordingly, the disclosure is intended to be illustrative, but not limiting, of the scope of the subject matter set forth in the following claims.

Claims

1. A computer-implemented method comprising: obtaining, using one or more processors, network telemetry data; identifying, using the one or more processors, a set of active machines based on network telemetry data; obtaining, using one or more processors, a first abandoned machine profile associated with a first active machine within the set of active machines, the first abandoned machine profile specific to a defined first subset of machines that includes the first active machine; obtaining, using one or more processors, a second abandoned machine profile associated with a second active machine within the set of active machines, the second abandoned machine profile specific to a defined second subset of machines that includes the second active machine; analyzing, using the one or more processors, a first subset of the network telemetry data associated with the first active machine based on the first abandoned machine profile; analyzing, using the one or more processors, a second subset of the network telemetry data associated with the second active machine based on the second abandoned machine profile; determining, using the one or more processors, one or more of the first active machine is an abandoned machine based on the first abandoned machine profile and the second active machine is an abandoned machine based on the second abandoned machine profile; and outputting, using the one or more processors, an identification of the one or more abandoned machines.

2. The method of claim 1 further comprising: generating a set of machines to monitor for abandonment based on a set of machines known to a network traffic monitor, and wherein the network telemetry data obtained is a subset of available network telemetry data available and is that network telemetry data that is associated with the set of machines to monitor for abandonment.

3. The method of claim 1 further comprising: determining a set of monitored machines, the set of monitored machines being monitored by a network traffic monitor that generates the network telemetry data, and the set of monitored machines is further being monitored for abandonment,

wherein the network telemetry data obtained describes the set of one or more monitored machines and includes the set of active machines.

4. The method of claim 3, wherein the set of monitored machines is identified by an identifier in the network telemetry data, wherein the identifier includes one or more of a range of a MAC addresses, a range of IP addresses, a list of MAC addresses, a list of IP addresses, a universal unique identifier of a virtual machine, and an interface ID of network adapter.

5. The method of claim 1, wherein analyzing the network telemetry data associated with the set of active machines determining whether the network telemetry data includes one or more of a first characteristic associated with an active but unabandoned machine and a second characteristic associated with an abandoned machine.

6. The method of claim 5, wherein the second characteristic associated with an abandoned machine is associated with one or more of a port number associated with traffic and a periodicity of the traffic, the traffic described by the network telemetry data.

7. The method of claim 1, wherein determining the one or more abandoned machines in the set of active machines includes: analyzing the network telemetry data; determining that the network telemetry data associated with a first machine is uncharacteristic of human use; and identifying the first machine as one of the one or more abandoned machines.

8. The method of claim 1, wherein the set of active machines includes one or more of a virtual machine and a physical machine.

9. The method of claim 1 further comprising: deactivating a first abandoned machine in the one or more identified, abandoned machines, wherein the deactivation is one of: automatic and based on the identification of the first abandoned machine as abandoned, or based on the identification of the first abandoned machine as abandoned and responsive to user input requesting abandonment.

10. The method of claim 1, wherein the identification of the one or more abandoned machines is output in a report presented to a user.

11. The method of claim 1, wherein the defined first subset of machines is associated with a class of machines automatically determined and associated with the first abandoned machine profile.

12. The method of claim 1, wherein the first subset of machines and the second subset of machines are associate with one or more of a different user set and a different environment.

13. A system comprising: one or more processors; and a memory storing instructions that, when executed by the one or more processors, cause the system to: obtain network telemetry data; identify a set of active machines based on network telemetry data; obtaining a first abandoned machine profile associated with a first active machine within the set of active machines, the first abandoned machine profile specific to a defined first subset of machines that includes the first active machine; obtaining a second abandoned machine profile associated with a second active machine within the set of active machines, the second abandoned machine profile specific to a defined second subset of machines that includes the second active machine; analyze a first subset of the network telemetry data associated with the first active machine based on the first abandoned machine profile; analyze a second subset of the network telemetry data associated with the second active machine based on the second abandoned machine profile; determine one or more of the first active machine is an abandoned machine based on the first abandoned machine profile and the second active machine is an abandoned machine based on the second abandoned machine profile; and output an identification of the one or more abandoned machines.

14. The system of claim 13 comprising instructions that, when executed, cause the system to: generate a set of machines to monitor for abandonment based on a set of machines known to a network traffic monitor, and wherein the network telemetry data obtained is a subset of available network telemetry data available and is that network telemetry data that is associated with the set of machines to monitor for abandonment.

15. The system of claim 13 comprising instructions that, when executed, cause the system to: determine a set of monitored machines, the set of monitored machines being monitored by a

network traffic monitor that generates the network telemetry data, and the set of monitored machines is further being monitored for abandonment, wherein the network telemetry data obtained describes the set of one or more monitored machines and includes the set of active machines.

16. The system of claim 15, wherein the set of monitored machines is identified by an identifier in the network telemetry data, wherein the identifier includes one or more of a range of a MAC addresses, a range of IP addresses, a list of MAC addresses, a list of IP addresses, a universal unique identifier of a virtual machine, and an interface ID of network adapter.

17. The system of claim 13, wherein analyzing the network telemetry data associated with the set of active machines determining whether the network telemetry data includes one or more of a first characteristic associated with an active but unabandoned machine and a second characteristic associated with an abandoned machine.

18. The system of claim 17, wherein the second characteristic associated with an abandoned machine is associated with one or more of a port number associated with traffic and a periodicity of the traffic, the traffic described by the network telemetry data.

19. The system of claim 13, wherein determining the one or more abandoned machines in the set of active machines includes: analyzing the network telemetry data; and determining that the network telemetry data associated with a first machine is uncharacteristic of human use; and identifying the first machine as one of the one or more abandoned machines.

20. The system of claim 13 comprising instructions that, when executed, cause the system to: deactivate a first abandoned machine in the one or more identified, abandoned machines, wherein the deactivation is one of: automatic and based on the identification of the first abandoned machine as abandoned, or based on the identification of the first abandoned machine as abandoned and responsive to user input requesting abandonment.
