

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12407649
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Connors, Jr.; Stephen Craig

Seamlessly verifying upgrades using mirrored dataplanes

Abstract

A system and method are provided for implementing a network component and verifying an update of the network component. The network component can be, e.g., a software-defined wide area network, a firewall, a router, or a load balancer. The network component can be an embedded network edge device that is implemented, e.g., in software, in circuitry, or using hardware acceleration (e.g., a data processing unit (DPU), a smart network interface card (SmartNIC), etc.). The updated version of the network component is verified by implementing it on a shadow dataplane concurrently with the current version operating on a primary dataplane, and comparing the performances of these two versions. Based on this comparison satisfying various verification criteria, the updated version passes a verification test and can be promoted to the primary dataplane.

Inventors:	Connors, Jr.; Stephen Craig (Nashville, TN)
Applicant:	Cisco Technology, Inc. (San Jose, CA)
Family ID:	94371437
Assignee:	Cisco Technology, Inc. (San Jose, CA)
Appl. No.:	18/410207
Filed:	January 11, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250039051 A1	Jan. 30, 2025

Related U.S. Application Data

us-provisional-application US 63516448 20230728

Publication Classification

Int. Cl.: **H04L9/40** (20220101); **G06F11/362** (20250101); **H04L9/32** (20060101); **H04L41/0806** (20220101); **H04L41/0816** (20220101); **H04L41/082** (20220101); **H04L41/0869** (20220101); **H04L41/0895** (20220101); **H04L43/50** (20220101); **H04L45/24** (20220101); **H04L45/80** (20220101)

U.S. Cl.:

CPC **H04L63/0236** (20130101); **G06F11/3636** (20130101); **H04L9/3247** (20130101); **H04L41/0806** (20130101); **H04L41/0816** (20130101); **H04L41/082** (20130101); **H04L41/0869** (20130101); **H04L41/0895** (20220501); **H04L43/50** (20130101); **H04L45/24** (20130101); **H04L45/80** (20220501); **H04L63/0218** (20130101); **H04L63/0263** (20130101); **H04L63/0435** (20130101); **H04L63/123** (20130101); **H04L63/1425** (20130101); **H04L63/1433** (20130101); **H04L63/145** (20130101); **H04L63/166** (20130101); **H04L63/20** (20130101);

Field of Classification Search

CPC: H04L (63/0236); H04L (9/3247); H04L (41/0806); H04L (41/0816); H04L (41/082); H04L (41/0869); H04L (41/0895); H04L (43/50); H04L (45/24); H04L (45/80); H04L (63/0218); H04L (63/0263); H04L (63/0435); H04L (63/123); H04L (63/1425); H04L (63/1433); H04L (63/145); H04L (63/166); H04L (63/20); H04L (41/0894); H04L (45/08); H04L (63/0272); G06F (11/3636)

USPC: 709/231

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9600263	12/2016	Stolarchuk et al.	N/A	N/A
2011/0188386	12/2010	Nicholson et al.	N/A	N/A
2016/0261460	12/2015	Shi et al.	N/A	N/A
2018/0241627	12/2017	Rodrigues	N/A	H04L 41/0813
2020/0236045	12/2019	Bays et al.	N/A	N/A
2020/0280566	12/2019	Raj et al.	N/A	N/A
2020/0313999	12/2019	Lee et al.	N/A	N/A
2023/0061613	12/2022	Cardozo et al.	N/A	N/A
2024/0267317	12/2023	Ameling	N/A	H04L 43/062
2025/0039051	12/2024	Connors	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2907025	12/2017	EP	N/A
3695565	12/2020	EP	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion for International Application No. PCT/US2025/011177, mailed Apr. 11, 2025, 15 Pages. cited by applicant

Primary Examiner: Tang; Karen C

Attorney, Agent or Firm: Polsinelli P.C.

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application claims priority to U.S. provisional application No. 63/516,448, titled “Data Processing Units (DPUs) and extended Berkley Packet Filters (eBPFs) for Improved Security,” and filed on Jul. 28, 2023, which is expressly incorporated by reference herein in its entirety.

BACKGROUND

(1) In software as a service (SaaS) deployments, upgrades to the software can be seamless with little to no impact on the users. This is achieved by performing the upgrades in a manner that is largely abstracted away from users, e.g., by rolling upgrades across multiple Kubernetes containers, by slowly shifting the load from the old version to the new versions (e.g., blue/green deployment), and monitoring the new versions. This allows continuous integration, continuous deployment (CI/CD) where the SaaS software can be kept updated.

(2) In contrast, for infrastructure and devices at the network edge (e.g. SD-WAN appliances, firewalls, and load balancers), upgrades and maintenance have been more disruptive to users. More particularly, upgrading embedded devices at the network edge has presented several challenges. First, these infrastructure upgrades often introduce downtime due to device failover and/or route re-convergence, and therefore these infrastructure upgrades can require a scheduled maintenance window. Second, these infrastructure upgrades often entail exhaustive pre- and post-upgrade

checks to ensure that the new software or policy does not negatively affect the network. Third, in case the upgrade fails, these infrastructure upgrades often include rollback and other contingency plans. Fourth, in-house testing, which occurs before the deployment/production phase, can fail to identify issues due to differences between the in-house settings/environment and the production settings/environment (e.g., the customer's own network). Thus, even after in-house testing of the upgrade, uncertainty remains because the in-house testing might fail to identify issues related to unique to characteristics of the customer's own network.

(3) Accordingly, improved methods and systems are desired for upgrading network edge devices. For example, improvements are desired that allow for seamless upgrades that are not disruptive to users of the network.

Description

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

(1) In order to describe the manner in which the above-recited and other advantages and features of the disclosure can be obtained, a more particular description of the principles briefly described above will be rendered by reference to specific embodiments thereof which are illustrated in the appended drawings. Understanding that these drawings depict only exemplary embodiments of the disclosure and are not therefore to be considered to be limiting of its scope, the principles herein are described and explained with additional specificity and detail through the use of the accompanying drawings in which:

(2) FIG. 1A illustrates a block diagram of a first example of a network, in accordance with certain embodiments.

(3) FIG. 1B illustrates a block diagram of a second example of the network, in accordance with certain embodiments.

(4) FIG. 2 illustrates a block diagram of a network device, in accordance with certain embodiments.

(5) FIG. 3 illustrates a flow diagram for an example of a method of implementing a software development lifecycle (SDLC), in accordance with certain embodiments.

(6) FIG. 4 illustrates a state diagram for an example of a step for deploying a new software/policy version using dual dataplanes, in accordance with certain embodiments.

(7) FIG. 5 illustrates a block diagram of a first example of implementing dual dataplanes in data processing units (DPUs), in accordance with certain embodiments.

(8) FIG. 6 illustrates a block diagram of a second example of implementing dual dataplanes in data processing units (DPUs), in accordance with certain embodiments.

(9) FIG. 7A illustrates a block diagram of an example of a configuration of pairwise primary and shadow dataplanes when there are eight total dataplanes, in accordance with certain embodiments.

(10) FIG. 7B illustrates a block diagram of an example of a scale-out configuration with eight dataplanes, in accordance with certain embodiments.

(11) FIG. 8A illustrates a flow diagram for an example of a method of training an artificial intelligence (AI) engine, in accordance with certain embodiments.

(12) FIG. 8B illustrates a flow diagram for an example of a method of using the trained AI engine, in accordance with certain embodiments.

DESCRIPTION OF EXAMPLE EMBODIMENTS

(13) Various embodiments of the disclosure are discussed in detail below. While specific implementations are discussed, it should be understood that this is done for illustration purposes only. A person skilled in the relevant art will recognize that other components and configurations may be used without parting from the spirit and scope of the disclosure.

Overview

(14) In one aspect, a method is provided for implementing a network component and verifying an

update to the network component. The method includes receiving, at one or more ports of a network device, ingress traffic comprising data packets; and processing a first subset of data packets from the ingress traffic at a first dataplane processes, the first subset being processed in accordance with first networking instructions to generate first egress data packets and a first record that represents a performance of the first dataplane when processing the first subset of data packets.

(15) The method further includes processing a second subset of data packets from the ingress traffic at a second dataplane processes, the second subset being processed in accordance with second networking instructions to generate second egress data packets and a second record that represents a performance of the second dataplane when processing the second subset of data packets, the second subset being mirrored to the first subset such that the second subset is identical to the first subset; and comparing, at a controller, the first record with the second record and using the comparison of the first record with the second record to verify an updated version of the network component. The first networking instructions execute a current version of the network component and the second networking instructions execute the updated version of the network component.

(16) In another aspect, the method may also include that the controller verifies the updated version of the network component by determining that the comparison of the first record with the second record passes a verification test based on the comparison of the first record with the second record satisfying one or more predefined criteria.

(17) In another aspect, the method may also include determining the predefined criteria based on verification metadata, the verification metadata providing indicia regarding predicted differences between the performance of the first dataplane with respect to the performance of the second dataplane.

(18) In another aspect, the method may also include that the verification metadata comprises learned values that are based on pre-verification testing the updated version of the network component in a test bed that simulates a commercial network, wherein the pre-verification testing occurs prior to the verification testing.

(19) In another aspect, the method may also include verifying the updated version of the network component while operating in a production environment in which the first egress data packets are used in a commercial network during the verification testing.

(20) In another aspect, the method may also include that the one or more predefined criteria of the verification test include that, for predefined parameters, values of the predefined parameters in the first record are within respective predefined ranges for the values of the predefined parameters in the second record, and the predefined parameters are selected from the group of performance parameters consisting of: (i) a minimum central processor unit (CPU) usage, (ii) a maximum CPU usage, (iii) an average CPU usage, (iv) a minimum memory usage, (v) a maximum memory usage, (vi) an average memory usage, (vii) a memory growth over a verification period, and (viii) a packet latency for packets to traverse a dataplane.

(21) In another aspect, the method may also include that the one or more predefined criteria of the verification test of comprise a criterion that a traffic volume of the first egress data packets is within a predefined range with respect to a traffic volume of the second egress data packets.

(22) In another aspect, the method may also include that the network component implements a network policy, and the updated version of the network component represents a modification of the network policy relative to the current version of the network component.

(23) In another aspect, the method may also include that the one or more predefined criteria of the verification test include: a first criterion that the updated version of the network policy does not adversely affect network traffic more than a predefined amount, and/or a second criterion that the updated version of the network policy provides an expected change to a predefined aspect of the network traffic.

(24) In another aspect, the method may also include receiving, at a packet dispatcher, the ingress traffic from the one or more ports; determining, by the packet dispatcher, which of the data packets

from the ingress traffic to include in the first subset and then transmitting the first subset to the first dataplane; and determining, by the packet dispatcher, which of the data packets from the ingress traffic to include in the second subset and then transmitting the second subset to the second dataplane.

(25) In another aspect, the method may also include that a first data processing unit (DPU) implements the first dataplane and a second DPU implements the second dataplane, and the packet dispatcher is implemented in only one of the first DPU and the second DPU, and the packet dispatcher spans data packets to the DPU of the first DPU and the second DPU in which the packet dispatcher is not implemented.

(26) In another aspect, the method may also include that the egress traffic transmitted by the one or more ports comprises the first egress traffic and excludes the second egress traffic.

(27) In another aspect, the method may also include that the first dataplane and the second dataplane are implemented in one or more data processing units (DPUs) using hardware acceleration to perform the first networking instructions and second networking instructions.

(28) In another aspect, the method may also include that the network component is configured to provide data-packet filtering, load balancing, security screening, malware detection, firewall protection, data-packet routing, data-packet switching, data-packet forwarding, computing header checksums, or implementing network policies.

(29) In another aspect, the method may also include that the first dataplane is configured to operate as a primary dataplane and the second dataplane is configured to operate as a shadow dataplane, such that the primary dataplane provides a functionality of the network component and the shadow dataplane is only used to verify the updated version of the network component.

(30) In another aspect, the method may also include storing state information in a memory that is accessible to the controller, the first dataplane, and the second dataplane, such that the first dataplane and the second dataplane are stateless.

(31) In another aspect, the method may also include that the apparatus is an edge-computing processor that is configured in an embedded device of a network edge.

(32) In another aspect, the method may also include that the apparatus is implemented as a virtual machine in a central processing unit (CPU), as a Berkeley packet filter (BPF).

(33) In one aspect, a computing apparatus includes a processor. The computing apparatus also includes a memory storing instructions that, when executed by the processor, configure the apparatus to perform the respective steps of any one of the aspects of the above recited methods.

(34) In one aspect, an apparatus is provided for implementing a network component and verifying an update to the network component. The apparatus includes one or more ports that receive ingress traffic and transmit egress traffic. The apparatus further includes circuitry comprising a first dataplane, a second dataplane, and a controller. The first dataplane processes, in accordance with first networking instructions, a first subset of data packets from the ingress traffic to generate first egress data packets and a first record representing a performance of the first dataplane when processing the first subset of data packets. The second dataplane processes, in accordance with second networking instructions, a second subset of data packets from the ingress traffic to generate second egress data packets and a second record representing a performance of the second dataplane when processing the second subset of data packets, the second subset being mirrored to the first subset such that the second subset is identical to the first subset. The controller receives and compares the first record with the second record and uses the comparison of the first record with the second record to verify an updated version of the network component. The first networking instructions execute a current version of the network component and the second networking instructions execute the updated version of the network component.

(35) In another aspect, the apparatus may also include that the controller verifies the updated version of the network component by determining that the comparison of the first record with the second record passes a verification test based on the comparison of the first record with the second

record satisfying one or more predefined criteria.

(36) In another aspect, the apparatus may also include that the circuitry is configured to determine the predefined criteria based on verification metadata, the verification metadata providing indicia regarding predicted differences between the performance of the first dataplane with respect to the performance of the second dataplane.

(37) In another aspect, the apparatus may also include that the verification metadata comprises learned values that are based on pre-verification testing the updated version of the network component in a test bed that simulates a commercial network, wherein the pre-verification testing occurs prior to the verification testing.

(38) In another aspect, the apparatus may also include that the circuitry is configured to verify the updated version of the network component while operating in a production environment in which the first egress data packets are used in a commercial network during the verification testing.

(39) In another aspect, the apparatus may also include that the one or more predefined criteria of the verification test include that, for predefined parameters, values of the predefined parameters in the first record are within respective predefined ranges for the values of the predefined parameters in the second record, and the predefined parameters are selected from the group of performance parameters consisting of: (i) a minimum central processor unit (CPU) usage, (ii) a maximum CPU usage, (iii) an average CPU usage, (iv) a minimum memory usage, (v) a maximum memory usage, (vi) an average memory usage, (vii) a memory growth over a verification period, and (viii) a packet latency for packets to traverse a dataplane.

(40) In another aspect, the apparatus may also include that the one or more predefined criteria of the verification test of comprise a criterion that a traffic volume of the first egress data packets is within a predefined range with respect to a traffic volume of the second egress data packets.

(41) In another aspect, the apparatus may also include that the network component implements a network policy, and the updated version of the network component represents a modification of the network policy relative to the current version of the network component.

(42) In another aspect, the apparatus may also include that the one or more predefined criteria of the verification test include: a first criterion that the updated version of the network policy does not adversely affect network traffic more than a predefined amount, and/or a second criterion that the updated version of the network policy provides an expected change to a predefined aspect of the network traffic.

(43) In another aspect, the apparatus may also include a packet dispatcher that: receives the ingress packets from the one or more ports, determines which of the data packets from the ingress traffic to include in the first subset and then transmit the first subset to the first dataplane, and determines which of the data packets from the ingress traffic to include in the second subset and then transmit the second subset to the second dataplane.

(44) In another aspect, the apparatus may also include that the circuitry comprises a first data processing unit (DPU) that implements the first dataplane and a second DPU that implements the second dataplane, and the packet dispatcher is implemented in only one of the first DPU and the second DPU, and the packet dispatcher spans data packets to the DPU of the first DPU and the second DPU in which the packet dispatcher is not implemented.

(45) In another aspect, the apparatus may also include that the egress traffic transmitted by the one or more ports comprises the first egress traffic and excludes the second egress traffic.

(46) In another aspect, the apparatus may also include that the circuitry comprises one or more data processing units (DPUs) that use hardware acceleration to perform the first networking instructions and second networking instructions.

(47) In another aspect, the apparatus may also include that the network component is configured to provide data-packet filtering, load balancing, security screening, malware detection, firewall protection, data-packet routing, data-packet switching, data-packet forwarding, computing header checksums, or implementing network policies.

- (48) In another aspect, the apparatus may also include that the first dataplane is configured to operate as a primary dataplane and the second dataplane is configured to operate as a shadow dataplane, such that the primary dataplane provides a functionality of the network component and the shadow dataplane is only used to verify the updated version of the network component.
- (49) In another aspect, the apparatus may also include a memory that stores state information, the memory being accessible to the controller, the first dataplane, and the second dataplane, such that the first dataplane and the second dataplane are stateless.
- (50) In another aspect, the apparatus may also include that the apparatus is an edge-computing processor that is configured in an embedded device of a network edge.
- (51) In another aspect, the apparatus may also include that the apparatus is implemented as a virtual machine in a central processing unit (CPU), as a Berkeley packet filter (BPF).

EXAMPLE EMBODIMENTS

- (52) Additional features and advantages of the disclosure will be set forth in the description which follows, and in part will be obvious from the description, or can be learned by practice of the herein disclosed principles. The features and advantages of the disclosure can be realized and obtained by means of the instruments and combinations particularly pointed out in the appended claims. These and other features of the disclosure will become more fully apparent from the following description and appended claims, or can be learned by the practice of the principles set forth herein.
- (53) The disclosed technology addresses the need in the art for improvements in upgrading network components. For example, a major challenge facing embedded devices at the network edge is the inability to seamlessly upgrade the embedded devices. These edge-computing devices can include a control plane that controls a dataplane in which data packets are received at various ports, interacted with in some manner (e.g., filtered, routed, forward, processed through a firewall, etc.), and then transmitted from the various ports, as discussed below.
- (54) Generally, upgrading an edge-computing device presents several challenges. First, downtime can result from upgrading the edge-computing device due to device failover and/or route re-convergence. Accordingly, a maintenance window can be scheduled, and the upgrade performed during the maintenance window to allow for the above-noted contingencies. Second, to ensure that the new software or policy does not have negative effects on the network, exhaustive pre-upgrade checks and post-upgrade checks can be performed on the edge-computing device or network component. Third, in case of upgrade failures, a rollback and other contingency plans can be used to rectify the upgrade failures. Fourth, the upgrade can be accompanied by uncertainty about issues in the new version. For example, uncertainty about issues in the new version might not have been identified in quality assurance (QA) checks. In some cases, for example, during the staging/testing phase, the testing environment used to initially verify the new version may be different than the customer's own network on which the new version is ultimately applied (e.g., the production/deployment phase). These differences may be due to unique characteristics of the customer's own network.
- (55) The systems and methods disclosed herein address the above-noted challenges by using dual dataplanes, including a primary dataplane and a shadow dataplane. For example, the primary dataplane executes a current version of the software or network policy, and the shadow dataplane executes a new version of the software or network policy. The shadow dataplane is used to perform verification testing of the new version by comparing its performance to that of the current version. Thus, the upgrade can undergo verification testing in the same environment as the current version is operating in (i.e., the customer's own network), thereby eliminating uncertainty about issues in the new version that may not have been identified in QA due to unique characteristics of the customer's own network.
- (56) Further, because the new version is verified in the shadow dataplane rather than the primary dataplane, the need for rollback and other contingency plans in case of upgrade failures can be largely mitigated. That is, until the verification testing is complete and the new version is promoted

to the primary dataplane, the current version continues to operate in parallel with the new version, and the network functionality continues to be performed by the current version rather than the new version. Then during promotion, which occurs after the new version passes verification testing, the new version can be gradually and gracefully transitioned to assuming the role of the new primary dataplane (i.e., the function of the network device is taken over by the new version). For example, if the new version fails the verification testing, there is no need to rollback to the current version because the current version is still operating to provide the functionality of the edge-computing device, unless and until the new version passes the verification testing. Further, the assurances provided by the pre-upgrade checks and post-upgrade checks can be (largely) integrated into the verification testing. Moreover, because the verification testing occurs in the background and is not disruptive to users, the upgrade can occur at any time rather than during a scheduled maintenance window.

(57) The systems and methods disclosed herein extend many of the advantages of seamless upgrades currently experienced for software as a service (SaaS) can be hereby experienced for infrastructure as a service (IaaS) in edge-computing devices and cloud computing environments. For example, in SaaS deployments, the above-noted challenges are largely abstracted away from users, especially the first and third of the above-noted challenges. In SaaS, this is achieved, e.g., by rolling upgrades across multiple containers. Further, in SaaS deployments, the above-noted challenges are largely abstracted away from users by slowly shifting the load from the old version to the new version (e.g., using blue/green deployment) and monitoring the new version. These strategies in SaaS deployments allow continuous integration, continuous deployment (CI/CD) where the SaaS software can be frequently and seamlessly updated in a manner that is invisible to the users.

(58) According to certain non-limiting examples, the systems and methods disclosed herein can achieve CI/CD for infrastructure, edge-computing components (e.g., hardware and software), and embedded edge devices, such that they can be frequently and seamlessly updated in a manner that is non-disruptive to the users (e.g., in ways that are different and/or similar to how this is achieved for SaaS). According to certain non-limiting examples, the systems and methods disclosed herein provide CI/CD in an embedded device that leverages artificial intelligence (AI) to mitigate the four challenges noted above.

(59) According to certain non-limiting examples, the systems and methods disclosed herein use the following four components: 1. A Control Plane Agent, which includes control logic representing the status of the dual dataplanes as well as the latest software and/or policy version(s) available (e.g., the current version and the new version); 2. A Primary Dataplane, which executes/implements the current stable version of software and/or policy; 3. A Shadow Dataplane, which executes/implements an upgrade or new version of software and/or policy; and 4. A Packet Dispatcher, which can provide multiple optional functionalities, including, e.g., a normal mode, a scale-out mode, a verification mode, and a promotion mode.

(60) According to certain non-limiting examples, the systems and methods disclosed herein are configured to operate using the following functionalities, which are performed by the system components: 1. Normal Mode—receives packets (e.g., ingress traffic) and then relays/sends the received packets to the primary dataplane, while the shadow dataplane is maintained in a passive/standby configuration; 2. Scale-out Mode—load balances packets across both dataplanes, e.g., by the packet dispatcher sending different subsets of the received packets to the primary dataplane and the shadow dataplane and using both the primary dataplane and the shadow dataplane to apply the current stable version to the different subsets of the received packets; 3. Verification Mode—receives packets and sends them to the primary dataplane, and additionally sends identical copies to the shadow dataplane, which applies the new version to these identical copies (e.g., in an embodiment that implements the respective dataplanes in data processing units (DPUs), the packet dispatcher may be part of the DPU itself and not a separate component); and 4.

Promotion Mode—once verification is complete, the shadow dataplane can take over as the primary dataplane, and a seamless transition occurs due to the graceful migration of data flows from the “Old Primary” to the “New Primary” dataplane.

(61) Examples of such edge-computing devices can include, but are not limited to, software-defined wide area network (SD-WAN) appliances, firewalls, load balancers, routers, switches, data processing units (DPUs), virtual machines that are implemented on one or more processors (e.g., a central processing unit (CPU)) for performing network functions or implementing network policies, or another component or device implemented at a network edge.

(62) According to certain non-limiting examples, the network edge device can include the following three planes: (i) the dataplane, which processes the transit traffic; (ii) the control plane, which sends and receives control signals to monitor and control the transit traffic; and (iii) the management plane, which interacts with the user or the network management system (NMS).

(63) Consider, for example, the operation of a router as an illustrative network edge device. Interfaces, IP subnets, and routing protocols can be configured through management plane protocols, including, e.g., a command-line interface (CLI), Network Configuration Protocol (NETCONF), and a northbound Representational State Transfer (REST) Application Programming Interface (API). The router runs control plane routing protocols (e.g., Open Shortest Path First (OSPF), Enhanced Interior Gateway Routing Protocol (EIGRP), Border Gateway Protocol (BGP), etc.) to discover adjacent devices and the overall network topology, or to discover reachability information in case of distance/path vector protocols). The router inserts the results of the control-plane protocols into Routing Information Base (RIB) and Forwarding Information Base (FIB). The dataplane software or ASICs, e.g., then use the FIB structures to forward the transit traffic. The management plane protocols (e.g., Simple Network Management Protocol (SNMP)) can then be used to monitor the device operation, its performance, interface counters, etc.

(64) Continuing with the non-limiting example of the network edge device being a router, in addition to controlling the routing protocols, the control plane protocols can also perform numerous other functions including: (i) interface state management (e.g., Point-to-Point Protocol (PPP), Transmission Control Protocol (TCP), and Link Aggregation Control Protocol (LACP)); (ii) connectivity management (e.g., Bidirectional Forwarding Detection (BFD), Connectivity Fault Management (CFM), etc.); (iii) adjacent device discovery (e.g., “hello” mechanisms present in most routing protocols, such as, End System-to-Intermediate System (ES-IS), Address Resolution Protocol (ARP), Internet Protocol version 6 (IPv6_Neighbor Discovery Protocol (NDP), Universal Plug and Play (UPnP) Simple Service Discovery Protocol (SSDP), etc.); (iv) topology or reachability information exchange (IP/IPv6 routing protocols, Intermediate System to Intermediate System (IS-IS) in Transparent Interconnection of Lots of Links (TRILL) and Shortest Path Bridging (SPB), Spanning Tree Protocol (STP), etc.); and (v) service provisioning (e.g., Resource Reservation Protocol (RSVP) for IntServ or Traffic Engineering (TE) based on Multiprotocol Label Switching (MPLS), uPNP SOAP (Simple Object Access Protocol) calls, etc.).

(65) Still continuing with the non-limiting example of the network edge device being a router, in addition to forwarding packets, the dataplane can also perform the following functions: (i) network address translation (NAT) session creation and NAT table maintenance; (ii) neighbor address gleaning (e.g., dynamic Media Access Control (MAC) address learning in bridging, IPv6 Source Address Validation Improvement (SAVI), etc.); (iii) NetFlow or sampled flow (sFlow) accounting; (iv) network access control list (ACL) logging; and (v) Error signaling, such as Internet Control Message Protocol (ICMP).

(66) According to certain non-limiting examples, the management and control planes can be implemented in a central processing unit (CPU) or in a data processing unit (DPU). According to certain non-limiting examples, the data plane could be implemented in numerous ways, including, e.g.: (i) as optimized code running on the same CPU as the control plane; (ii) as code running on a dedicated CPU core (e.g., a dedicated CPU for high-speed packet switching, such as a Linux

server); (iii) as code running on linecard CPUs (e.g., a CISCO 7200 series router); (iv) as dedicated processors (e.g., network process units (NPUs), data process units (DPUs), smart network interface cards (SmartNICs), etc.); and (v) as switching hardware (application-specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), etc.); and (vi) as switching hardware on numerous linecards.

(67) According to certain non-limiting examples, the dataplane receives and processes the ingress packets. Further, the dataplane can selectively forward packets destined for the router (e.g., Secure Shell (SSH) traffic or routing protocol updates) or packets that need special processing (e.g., IP datagrams with IP options or IP datagrams that have exceeded their TTL) to the control plane.

(68) According to certain non-limiting examples, the management ports on some devices (e.g. data center switches) can be connected directly to a control-plane CPU and thus bypass a switching ASIC.

(69) According to certain non-limiting examples, the control plane can pass outbound packets to the dataplane, or use its own forwarding mechanisms to determine the outgoing interface and the next-hop router (e.g., when using the local policy routing).

(70) A computer network is a geographically distributed collection of nodes interconnected by communication links and segments for transporting data between end nodes, such as personal computers, cellular phones, workstations, or other devices, such as sensors, etc. Many types of networks are available, with the types ranging from local area networks (LANs) to wide area networks (WANs). LANs typically connect the nodes over dedicated private communications links located in the same general physical location, such as a building or campus. WANs, on the other hand, typically connect geographically dispersed nodes over long-distance communications links, such as common carrier telephone lines, optical light paths, synchronous optical networks (SONET), or synchronous digital hierarchy (SDH) links, or Powerline Communications (PLC) such as IEEE 61334, IEEE P1901.2, and others. The Internet is an example of a WAN that connects disparate networks throughout the world, providing global communication between nodes on various networks. The nodes typically communicate over the network by exchanging discrete frames or packets of data according to predefined protocols, such as the Transmission Control Protocol/Internet Protocol (TCP/IP). In this context, a protocol consists of a set of rules defining how the nodes interact with each other. Computer networks may be further interconnected by an intermediate network node, such as a router, to forward data from one network to another.

(71) FIG. 1A is a schematic block diagram of a non-limiting example of a computer network **100** that includes various nodes/devices, such as a plurality of routers/devices interconnected by links or networks. For example, customer edge (CE) routers **110** may be interconnected with provider edge (PE) routers (e.g., PE-1 **120a**, PE-2 **120b**, and PE-3 **120c**) in order to communicate across a core network, such as an illustrative network backbone **130**. For example, the routers (e.g., CEs **110**, PEs **120**) may be interconnected by the public Internet, a multiprotocol label switching (MPLS), or a virtual private network (VPN). Data packets **140** (e.g., traffic/messages) may be exchanged among the nodes/devices of the computer network **100** over links using predefined network communication protocols such as the Transmission Control Protocol/Internet Protocol (TCP/IP), User Datagram Protocol (UDP), Asynchronous Transfer Mode (ATM) protocol, Frame Relay protocol, or any other suitable protocol. Those skilled in the art will understand that any number of nodes, devices, links, etc. may be used in the computer network, and that the view shown herein is for simplicity.

(72) In some implementations, a router or a set of routers may be connected to a private network (e.g., dedicated leased lines, an optical network, etc.) or a virtual private network (VPN), such as an MPLS VPN utilizing a Service Provider network, via one or more links exhibiting very different network and service level agreement characteristics.

(73) According to certain non-limiting examples, a given customer site may fall under any of the following categories: 1.) Site Type A: a site connected to the network (e.g., via a private or VPN

link) using a single CE router and a single link, with potentially a backup link (e.g., a 3G/4G/5G/LTE backup connection). For example, a particular CE router **110** shown in network **100** may support a given customer site, potentially also with a backup link, such as a wireless connection. 2.) Site Type B: a site connected to the network using two MPLS VPN links (e.g., from different Service Providers) using a single CE router, with potentially a backup link (e.g., a 3G/4G/5G/LTE connection). A site of type B may itself be of different types: 2a.) Site Type B1: a site connected to the network using two MPLS VPN links (e.g., from different Service Providers), with potentially a backup link (e.g., a 3G/4G/5G/LTE connection). 2b.) Site Type B2: a site connected to the network using one MPLS VPN link and one link connected to the public Internet, with potentially a backup link (e.g., a 3G/4G/5G/LTE connection). For example, a particular customer site may be connected to network **100** via PE-3 and via a separate Internet connection, potentially also with a wireless backup link. 2c.) Site Type B3: a site connected to the network using two links connected to the public Internet, with potentially a backup link (e.g., a 3G/4G/5G/LTE connection). 3.) Site Type C: a site of type B (e.g., types B1, B2 or B3) but with more than one CE router (e.g., a first CE router connected to one link while a second CE router is connected to the other link), and potentially a backup link (e.g., a wireless 3G/4G/5G/LTE backup link). For example, a particular customer site may include a first CE router **110** connected to PE-2 **120b** and a second CE router **110** connected to PE-3 **120c**.

(74) FIG. 1B illustrates an example of network **100** in greater detail, according to various embodiments. As shown, network backbone **130** may provide connectivity between devices located in different geographical areas and/or different types of local networks. For example, network **100** may comprise local/branch networks **160** and **162** that include devices/nodes **172**, **174**, **176**, and **178** and devices/nodes **164** and **166**, respectively, as well as a data center/cloud environment **150** that includes servers **168** and **170**. Notably, local network **160**, local network **162**, and data center **150** can be located in different geographic locations.

(75) Server **168** and server **170** can include, in various embodiments, a network management server (NMS), a dynamic host configuration protocol (DHCP) server, a constrained application protocol (CoAP) server, an outage management system (OMS), an application policy infrastructure controller (APIC), an application server, etc. As would be appreciated, network **100** may include any number of local networks, data centers, cloud environments, devices/nodes, servers, etc.

(76) In some embodiments, the techniques herein may be applied to other network topologies and configurations. For example, the techniques herein may be applied to peering points with high-speed links, data centers, etc.

(77) In various embodiments, network **100** may include one or more mesh networks, such as an Internet of Things network. Loosely, the term “Internet of Things” or “IoT” refers to uniquely identifiable objects (things) and their virtual representations in a network-based architecture. In particular, the next frontier in the evolution of the Internet is the ability to connect more than just computers and communications devices, but rather the ability to connect “objects” in general, such as lights, appliances, vehicles, heating, ventilating, and air-conditioning (HVAC), windows and window shades and blinds, doors, locks, etc. The “Internet of Things” thus generally refers to the interconnection of objects (e.g., smart objects), such as sensors and actuators, over a computer network (e.g., via IP), which may be the public Internet or a private network.

(78) Notably, shared-media mesh networks, such as wireless or PLC networks, etc., are often deployed on what are referred to as Low-Power and Lossy Networks (LLNs), which are a class of network in which both the routers and their interconnect are constrained: LLN routers typically operate with constraints, e.g., processing power, memory, and/or energy (battery), and their interconnects are characterized by, illustratively, high loss rates, low data rates, and/or instability. LLNs are comprised of anything from a few dozen to thousands or even millions of LLN routers, and support point-to-point traffic (between devices inside the LLN), point-to-multipoint traffic (from a central control point such as the root node to a subset of devices inside the LLN), and

multi-point-to-point traffic (from devices inside the LLN towards a central control point). Often, an IoT network is implemented with an LLN-like architecture. For example, as shown, local network **160** may be an LLN in which CE-2 operates as a root node for devices/nodes 10-16 in the local mesh, in some embodiments.

(79) In contrast to traditional networks, LLNs face a number of communication challenges. First, LLNs communicate over a physical medium that is strongly affected by environmental conditions that change over time. Some examples include temporal changes in interference (e.g., other wireless networks or electrical appliances), physical obstructions (e.g., doors opening/closing, seasonal changes such as the foliage density of trees, etc.), and propagation characteristics of the physical media (e.g., temperature or humidity changes, etc.). The time scales of such temporal changes can range between milliseconds (e.g., transmissions from other transceivers) to months (e.g., seasonal changes of an outdoor environment). In addition, LLN devices typically use low-cost and low-power designs that limit the capabilities of their transceivers. In particular, LLN transceivers typically provide low throughput. Furthermore, LLN transceivers typically support limited link margin, making the effects of interference and environmental changes visible to link and network protocols. The high number of nodes in LLNs in comparison to traditional networks also makes routing, quality of service (QoS), security, network management, and traffic engineering extremely challenging, to mention a few.

(80) FIG. 2 is a schematic block diagram of an example node/device **200** (e.g., an apparatus) that may be used with one or more embodiments described herein, e.g., as any of the computing devices shown in FIG. 1A and FIG. 1B, the device **200** may also be any other suitable type of device depending upon the type of network architecture in place, such as IoT nodes, etc. Device **200** can include one or more edge devices **206**, one or more processor(s) **214**, and a memory **202** interconnected by a system bus **204**.

(81) The edge device **206** can include the mechanical, electrical, and signaling circuitry for communicating data over physical links coupled to the network **100**. The edge device **206** can be configured to transmit and/or receive data using a variety of different communication protocols. the edge device **206** can also be used to implement one or more virtual network interfaces, such as for virtual private network (VPN) access, known to those skilled in the art. The edge device **206** can be implemented as software executed on a central processing unit (CPU), such as a virtual machine like a Berkley packet filter (BPF) or extended BPF (eBPF) that is configured to implement a network policy, for example. Alternatively or additionally, The edge device **206** can be implemented as a separate piece of hardware (e.g., a data processing unit (DPU), a graphics processing unit (GPU), a smart network interface card (smartNIC), a network interface controller, an application-specific integrated circuit (ASIC), field programmable gate array (FPGA), or other device/circuitry configured to perform the function of a network component).

(82) The edge device **206** can be configured to provide one or more network functions, including, e.g., data-packet filtering, load balancing, packet screening, pattern detection for cybersecurity threats, malware detection, firewall protection, data-packet routing, data-packet switching, data-packet forwarding, computing header checksums, or implementing network policies. According to certain non-limiting examples, the edge device **206** can be an embedded device at the network edge. The edge device **206** can include (or be part of) a software-defined wide area network (SD-WAN) appliance, a firewall, or a load balancer, for example. Moreover, the systems and methods disclosed herein can be used with any edge device **206** that includes a dataplane that can be intermittently updated to a new version.

(83) The edge device **206** can include a dataplane, a control plane, and a management plane, as discussed above. Further, instructions implementing the control plane and the management plane can be stored and/or executed in the processor(s) **214**. Additionally or alternatively, the edge device **206** can include processors or circuits that implement one or more functions of the control plane and the management plane. The edge device **206** can include a series of ports (e.g., port 1 **226a**,

port 2 **226b**, port 3 **226c**, port 4 **226d**, and port 5 **226e**). The edge device **206** can also include a control agent **218**, a dispatcher **220**, a dataplane 1 **222**, and a dataplane 2 **224**.

(84) Memory **202** can include a plurality of storage locations that are addressable by the processor(s) **214** and the edge device **206** for storing software programs and data structures associated with the embodiments described herein. Memory **202** can include various data structures **208** and can include instructions for executing the operating system **210**, the current network component **212**, and the updated network component **216**. The processor(s) **214** can include logic adapted to execute the software programs and manipulate the data structures **208**. An operating system **210** (e.g., the Internetworking Operating System, or IOS®, of Cisco Systems, Inc., another operating system, etc.), portions of which can be in memory **202** and executed by the processor(s), functionally organizes the node by, inter alia, invoking network operations in support of software processors and/or services executing on the device.

(85) It will be apparent to those skilled in the art that other processor and memory types, including various computer-readable media, may be used to store and execute program instructions pertaining to the techniques described herein. Also, while the description illustrates various processes, it is expressly contemplated that various processes may be embodied as modules configured to operate in accordance with the techniques herein (e.g., according to the functionality of a similar process). Further, while processes may be shown and/or described separately, those skilled in the art will appreciate that processes may be routines or modules within other processes.

(86) FIG. **3** illustrates an example of method **300** for a software development life cycle (SDLC). Although the example method **300** depicts a particular sequence of operations, the sequence may be altered without departing from the scope of the present disclosure. For example, some of the operations depicted may be performed in parallel or in a different sequence that does not materially affect the function of the method **300**. In other examples, different components of an example device or system that implements method **300** may perform functions at substantially the same time or in a specific sequence.

(87) Method **300** provides a structured process to enable high-quality, low-cost dataplane development in a short period. Method **300** can produce a dataplane infrastructure that meets customer expectations with minimal interruption.

(88) According to some examples, in step **302**, the method includes planning new dataplane infrastructure.

(89) According to certain non-limiting examples, step **302** includes planning and requirement analysis. Requirement analysis can be performed by the senior members of the team with inputs from the customer, the sales department, market surveys, and domain experts in the industry. This information can be used to plan the basic project approach based on studies in the economical, operational, and technical areas.

(90) According to some examples, in step **304**, the method includes defining the dataplane infrastructure at step **304**.

(91) According to certain non-limiting examples, upon completion of the requirement analysis, the product requirements can be defined and documented. Further, these product requirements can be approved by the customer or the market analysts. This can be done using a software requirement specification (SRS) document which consists of all the product requirements to be designed and developed during the project life cycle.

(92) According to some examples, in step **306**, the method includes designing the dataplane infrastructure at step **306**.

(93) According to certain non-limiting examples, the SRS is used as the reference for product architects to come out with the best architecture for the product to be developed. Based on the requirements specified in SRS. According to certain non-limiting examples, more than one design approach can be proposed for the product architecture is proposed and documented in a DDS—Design Document Specification.

- (94) This DDS is reviewed by various stakeholders and a preferred design approach is selected based on various selection criteria (e.g., based on various parameters such as risk assessment, product robustness, design modularity, budget, and time constraints). A design approach defines the architectural modules of the product along with its communication and data flow representation with the external and third-party modules (if any).
- (95) According to some examples, in step **308**, the method includes building/developing the dataplane infrastructure at step **308**.
- (96) According to certain non-limiting examples, in step **308** the development is performed and the product is built. The programming code is generated as per DDS during this step. Developers follow the coding guidelines defined by their organization and programming tools like compilers, interpreters, debuggers, etc. are used to generate the code. Different high-level programming languages such as C, C++, Pascal, Java and PHP are used for coding.
- (97) According to some examples, in step **310**, the method includes testing the dataplane infrastructure at step **310**.
- (98) According to certain non-limiting examples, step **310** can include parts of all stages of the SDLC cycle, and can thus be viewed as a subset of all the stages of the SDLC model. Further, some testing activities can be integrated with other stages of the SDLC method **300**. Once a code commit is ready, testing can proceed by provisioning the code commit in a staging environment that is intended to be representative of how the new version will be used in practice. Then, testing proceeds by measuring various signals to determine that the product/new version functions as desired. Generally, step **310** can include testing the product/new version for defects, bugs, or security vulnerabilities, and then reporting, tracking, fixing, and retesting the defects, bugs, or security vulnerabilities, until the product reaches the quality standards and passes a quality assurance (QA) process.
- (99) According to some examples, in process **312**, the method includes deploying the dataplane infrastructure at process **312**. Process **312** can include step **314** and step **316**. In process **312**, the new version/product can be deployed by provisioning it in a production environment (e.g., the customer's own network) and performing additional testing in this environment to ensure that the new version/product by measuring various signals to determine that the product/new version functions as desired.
- (100) According to some examples, in step **314**, the method includes verifying new dataplane infrastructure operating on the shadow data plane at step **314**. According to certain non-limiting examples, step **314** can be performed using the verification mode **406** and the decision block **408**, which are illustrated in FIG. 4 and are described below with reference to FIG. 4.
- (101) According to some examples, in step **316**, the method includes promoting new dataplane infrastructure to primary dataplane infrastructure (e.g., by transitioning the data flow from the previous primary dataplane to the previous shadow dataplane, which becomes the new primary dataplane). According to certain non-limiting examples, step **316** can be performed using the promotion mode **410**, which is illustrated in FIG. 4 and is described below with reference to FIG. 4.
- (102) According to certain non-limiting examples, the verification process that is disclosed herein using the shadow dataplane can replace previous methods of verification testing. Alternatively, the verification process that is disclosed herein can be combined with, rather than replace, previous methods of verification testing (e.g., step **310**) that exist as a normal course of the Software Development Lifecycle (SDLC). For example, code commits can be continuously integrated and tested in a staging environment, prior to the use of the shadow-dataplane-based verification process in a production environment. Once staging-phase testing is complete, the code can merged into the production branch and pushed into production. For example, the shadow-dataplane-based verification process can provide a unique non-disruptive way as the last stage of the continuous integration, continuous deployment (CI/CD) promotion pipeline.
- (103) Failures (e.g., defects, bugs, or vulnerabilities) detected in the production/deployment testing

(i.e., during the Verification Mode) can indicate a field escape of a bug that was missed during the QA process (e.g., the testing in step **310**). Thus, the production/deployment testing (i.e., during the Verification Mode) provides a way to prevent the bug from negatively impacting production by doing out-of-band verification within the actual production environment (e.g., the customer's own network), but in a manner that is not disruptive and/or is invisible to the customer (e.g., without the disruption of software rollbacks).

(104) Failures (e.g., defects, bugs, or vulnerabilities) detected in the production/deployment testing can be treated as field escapes. Although verification-mode failures do not negatively impact production, they can, however, incur other costs associated with QA field escapes (e.g., delay in getting new code to production, need to reverse engineer the failure, need to augment QA testing in staging to ensure this type of escape does not reoccur). Thus, the Verification Mode failures can be treated as field escapes by automatically generating bug reports for the team to address in both the software-development phase and staging-test phase (e.g., step **308** and step **310**).

(105) FIG. **4** illustrates a non-limiting example of process **312**. According to certain non-limiting examples, step **302** goes from start **402** to the normal mode **404**.

(106) In normal mode **404**, packets are simply sent through the primary dataplane and processed as normal. The shadow dataplane is in standby, waiting in case a new software and/or policy that needs verification arrives. Generally, there is no way to predict when the new software and/or policy that needs verification will arrive.

(107) In scale-out mode **412**, the idle resources in the shadow dataplane can be put to use for load balancing. In the absence of a shadow dataplane, a mechanism for scale-out of the network edge device is to leverage resources along the application path (e.g., DPU in the server+DPU in the switch). However, there may be cases where a device is overloaded and there is no way to shift load around in the immediate term. When this is the case, additional resources may become temporarily available by switching to scale-out mode **412** to take advantage of the shadow dataplane. Scale-out mode **412** is initiated by a scale-out request **426**, and then terminated when there is no longer a need for additional resources, resulting in signal **424** indicating that the normal capacity threshold is maintained.

(108) According to certain non-limiting examples, in scale-out mode **412**, any verification testing currently being processed is canceled, and a pause is put on accepting verification testing. The system engages both dataplanes as Active/Active (i.e., the primary dataplane is in an active state and the shadow dataplane is in an active state, rather than standby). Both dataplanes implement a current version of the software/policy. The packet dispatcher load balances packets across the active dataplanes to use all available capacity. For example, in a pure x86 or Advanced RISC Machine (ARM) system, this can be implemented as another pipeline to process packets, spreading the workload across more cores. In a system with DPUs, e.g., this means taking advantage of the additional hardware capacity on the DPUs.

(109) According to certain non-limiting examples, in verification mode **406**, packets are mirrored through two versions of the software and/or policy, which are running concurrently. The primary dataplane can be agnostic to the fact that the system is in verification mode and simply functions as it normally would. The shadow dataplane implements the new software and/or policy, and the shadow dataplane drops packets at the end of the pipeline rather than transmitting them. That is, both the primary dataplane and the shadow dataplane receive the exact same packets, making it possible to perform a direct comparison between the performance of the current and new versions of the software/policy. Further, this comparison is performed using the customer's actual traffic, and this comparison is performed transparently, without any impact on the execution of the network.

(110) As part of (or separate from) the verification mode **406**, the decision block **408** includes checks to determine, after a specified testing interval, whether the new version passed or failed the verification test(s). If the verification fails, the fail signal **418** indicates that the shadow dataplane has not been authorized to take over as the primary dataplane. The system reverts to normal mode

404 without passing through promotion mode **410**, and the system sends a message to the controller indicating the cause of the failure. A newer policy or software version is then developed and the verification mode **406** is reinitiated using the newer policy or software version. If verification testing passes, the system generates a signal of pass **420** and moves to promotion mode **410**.

(111) According to certain non-limiting examples, the decision block **408** can determine whether the new version passes the verification test based on an accumulated confidence over a period of time. For example, a confidence score can increase as packets are processed through the primary dataplane and the shadow dataplane and the accumulated statistics and measurements/comparisons between the respective dataplanes fall within desired parameters. As this continues over a longer period of time without deviations outside of predefined limits, the confidence increases that the new version is behaving in an acceptable/desirable manner. When the confidence score exceeds a predefined value over a predefined testing period, then the new version is determined to have passed the verification test.

(112) According to certain non-limiting examples, after the predefined testing period, a confidence score below the predefined value indicates failure and corrective action is taken to modify the new version before subsequent verification testing. Alternatively, a failed verification test occurs when the confidence score is less than a failure threshold, a passed verification test occurs when the confidence score exceeds a pass threshold, and a confidence score between the failure threshold and pass threshold is inclusive, resulting in further verification testing before the verification testing generates a conclusive pass or fail result.

(113) According to certain non-limiting examples, after the predefined testing period, the results can be inconclusive and the testing period can be extended for another predefined period, after which either the confidence score exceeds the threshold and the new version passes or the new version fails.

(114) According to certain non-limiting examples, the confidence score is not required to reach 100% confidence. For example, reaching 100% confidence can depend on exercising/testing every single policy and feature of the new version, but often the customer can use only a subset of the available policies and features. Accordingly, a user can select a confidence level that satisfies the verification testing.

(115) According to certain non-limiting examples, the confidence space can be multi-dimensional, such that multiple confidence scores can be generated with respect to different dimensions of the space. For example, one dimension can relate to CPU usage, another dimension can relate to packet latency through the edge device **206**. A third dimension can relate to memory usage, and so forth.

(116) Once verification mode **406** is complete and the new version has passed the verification, promotion mode **410** is initiated to promote the shadow dataplane to the primary dataplane. Preferably, this promotion is performed gracefully to minimize disruption to the running system. Further, this promotion can be performed gradually, in case a rollback becomes required due to there being any issues that were missed during verification.

(117) According to certain non-limiting examples, in promotion mode **410**, the shadow dataplane becomes the “primary pending” dataplane and the packet dispatcher begins forwarding all new flows to the primary pending dataplane. Once these flows are observed to flow normally, existing flows are gracefully migrated from the primary dataplane to the primary pending dataplane until all flows are transiting through the primary pending dataplane. At this point, the primary dataplane moves to the shadow dataplane, and the primary pending dataplane to the primary dataplane. Then the signal for the primary and shadow role exchange **422** is triggered, sending the state machine representing process **312** back to normal mode **404**.

(118) Returning to verification mode **406**, verification can be predicated on assessing the difference in behaviors of two versions of network software and/or network policies. In many cases, a comparison between the two versions is more involved than merely observing that the performance

of the new version matches the performance of the current version. For example, the new version can be expected to perform differently than the current version. For example, the new version might be intended to improve performance in certain aspects relative to the current version. Thus, the comparison between the performances of the two versions can be informed regarding expected differences in performance due to improvements integrated into the new versions, and, in anticipation of these differences, the comparison can set out criteria for the verification testing that account for the expected differences in performance.

(119) In one case that exemplifies a simple verification, the new version of the software is a minor revision to the current software that is limited to optimizing the pipeline to improve performance. In this case, verification can be straightforward because the expected result of the new version of the software is that the same number of packets are received and transmitted, the same policies are matched, the CPU and memory usage are the same or lower, the latency of packets transiting the dataplane is reduced and the effective throughput increased. Thus, basic heuristics can be used to determine whether the new version is an improvement relative to the current version, which would signify a successful verification.

(120) In other cases, however, verification is not as simple as in the above example. In such cases, new versions of software and policy can be accompanied by metadata that can inform the system how to interpret the results of the comparison between performances of the current and new versions to establish predefined criteria for verification.

(121) In these more complicated/non-trivial verification cases, verification criteria can be established for various aspects of the performance comparison. For example, to determine that the new version of the software will not disrupt the network, various different factors can be considered for the metrification of the performance when executing a firewall dataplane, including, e.g., CPU performance, memory usage, packet latency, and traffic volume. For example, the predefined verification criteria can include metrics related to changes in the CPU usage, including, e.g., the minimum CPU usage, the maximum CPU usage, and the average CPU usage. Further, the predefined verification criteria can include metrics related to changes in memory usage, including, e.g., the minimum memory usage, the maximum memory usage, the average memory usage, and memory growth over the verification period. Additionally, the predefined verification criteria can include metrics related to changes in the packet latency, including, e.g., the average time it takes packets to traverse the dataplane.

(122) In addition to the above performance metrics, the predefined verification criteria can include various traffic-volume metrics. These traffic-volume metrics can include, e.g., the total number of packets processed, the number of dropped packets, and the number of packets that are transmitted (i.e., transmitted from the ports to other network devices). In certain verification cases, it is anticipated that traffic volume should be identical at egress, between the primary and shadow dataplanes. For example, if N packets arrived and K packets were dropped due to policy, L packets should be transmitted at the end. For verification to be precise, the exact same number of packets (e.g., N packets) can be ensured on both dataplanes. The exact same number of packets can be ensured, e.g., by sending an inline control packet that signals the start and end of verification. This ensures that both dataplanes are operating on the same N .

(123) Whereas traffic volume should be identical at egress for cases in which the policies or other aspects of the processing do not change the number of dropped packets, in other cases the number of dropped packets can change, but often the number of dropped packets will change in predictable ways, which can be communicated via the accompanying metadata to generate traffic-volume metrics that are indicative of the predicted changes. That is, some versions of software may alter the number of dropped packets (e.g., above K value) for valid reasons.

(124) For instance, in the case that the current version of the software does not properly enforce dropping packets with a specific Internet Control Message Protocol (ICMP) code, but the new version does properly enforce dropping packets with the specific ICMP code. In this case, the new

version can be expected to drop more packets than the current version. The percentage change in the number of dropped packets can be learned during the staging phase, and then this learned value can be included in the metadata. When comparing the performances of the primary and shadow dataplanes during the verification mode **406**, which occurs during the production phase of the SDLC, the metadata can be used to predict a range of acceptable values for the number of dropped packets for the shadow dataplane relative to the primary dataplane, and this predicted range can be one of the verification criteria.

(125) The use of metadata to predict values for verification criteria is not limited to traffic volume, as would be understood by a person of ordinary skill in the art. Generally, verification metadata can be used as part of the upgrade to communicate/train the system about what is going to be different in the new version of the software/policy. For example, the verification metadata indicates what is expected to be different in the new version of the software/policy, such that, in verification mode **406**, the expected differences are not treated as verification failures. That is, the verification metadata can inform the decision regarding whether the new version of the software/policy is good (i.e., passes verification testing). The verification metadata can be generated in various ways, including, e.g., by being hard coded by a network operations center (NOC) engineer, a programmer, or a software developer; generated during in-house testing or staging-phase testing; based on product specification documents or expected changes, etc. The system can then use the verification metadata to predict how the new version should perform relative to the current version in order to pass verification testing. For example, the system can include an artificial intelligence (AI) engine that has been trained to predict these differences between the respective performances of the current and new versions of the software/policy.

(126) For example, the AI engine can effectively replace the NOC engineer (i.e., the person who would typically perform the firewall upgrade) for production testing by implementing the verification criteria for verification mode **406** based on the verification metadata.

(127) For example, in the SDLC, a NOC engineer could perform the pre-production testing, push the new version of the software, and then do the post-upgrade testing. The NOC engineer might use a spreadsheet listing various performance values for the current version and the new version, where these values can include, e.g., the CPU usage, memory usage, and the number of packets. The spreadsheet might further list changes in policies, and so forth, and, then by comparing these sets of values for current and new versions, the NOC engineer applies their judgment two to make sure the values for the current and new versions are either the same or the differences are expected/justified.

(128) The AI engine can essentially replace the NOC engineer and can automate the above process by using the verification metadata to predict what differences are expected/justified. The AI engine can be integrated with the control-plane agent **508**, and the AI engine can be trained to predict what metrics the new version should meet to demonstrate that the new version is operating properly and therefore passes the verification testing. For example, there can be different types of new versions, such as upgrades to software, new policies, etc. And different types of dataplane upgrades (i.e., new versions) can have different influences on the performance.

(129) For example, some dataplane upgrades can be intended to decrease CPU usage without changing packet routing. In this case, the AI engine can predict that the new version is operating properly when the egress data packets from the respective dataplanes are identical and the CPU usage of the shadow dataplane is within a predefined range with respect to the CPU usage of the primary dataplane (e.g., the CPU usage of the shadow dataplane is between 10% and 30% less than the CPU usage of the primary dataplane or the CPU usage of the shadow dataplane is between 15% and 25% less than the CPU usage of the primary dataplane). A hashing algorithm can be used to efficiently compare the egress packets of the respective dataplanes and ensure that they are identical. For example, in each dataplane, a hash can be generated for each of the egress data packets, and the hashes can then be combined (e.g., by summing the hashes). Then the resulting combination (e.g., the sum of the hashes) from one dataplane can be compared to that of the other

dataplane. When the egress data packets are identical the combined hashes for the respective dataplanes will also be identical.

(130) The ranges and parameters used to compare the performance of the primary dataplane and the shadow dataplane can be based in part on verification metadata that is generated during the staging phase of testing that is prior to the provisioning phase of testing. As discussed above, the staging phase of testing can be performed by monitoring the new version in a controlled environment, whereas the provisioning phase of testing is realized in an actual production environment using the customer's actual data and using the dual dataplanes by deploying the new version in parallel with the current version. That is, code commits can be tested in a staging environment, prior to the use of the shadow-dataplane-based verification process in a production environment. Once staging-phase testing is complete, the code can be merged into the production branch and pushed into production.

(131) In the staging phase of testing, the new version undergoes functional quality assurance in a controlled environment, and observations and data generated during the staging phase of testing can be used to generate the verification metadata, and the verification metadata can be used by the AI engine to predict the parameters/ranges relating the performance of primary dataplane to the performance of the shadow dataplane, when the new version is operating correctly. Additionally or alternatively, the observations and data generated during the staging phase of testing can be used to train the AI engine. For example, the controlled environment can be used to process the same curated data sets using the current version and using the new version, and the respective performances can be compared for the versions to determine the ranges of performance values that would satisfy the verification testing.

(132) In addition to the above discussion, policy verification can also include some unique aspects with respect to the verification criteria. For example, policy verification can be distinct from software verification because, in many cases, the version of the policies is expected to change the behavior of the system. Thus, the verification criteria for a new version of the policies can include metrics quantifying/indicating whether the new policy would negatively impact user traffic and whether the new policy effectively delivers the intended outcome. For example, the verification criteria can include one or more metrics quantifying whether the new policy has exposed some deficiency in the system that would negatively impact user traffic (e.g. performance degradation due to high CPU usage). Additionally, the verification criteria can include one or more metrics that assist the administrator in determining whether the policy effectively delivers the outcome that the administrator instituted the new policy to achieve.

(133) FIG. 5 illustrates a non-limiting example of a dual dataplane architecture **500** for a network appliance. Another dual dataplane architecture for use with DPUs (i.e., the DPU dual dataplane architecture **600**) is presented in FIG. 6. In a physical or virtual machine that does not have access to DPUs, the systems and methods disclosed herein provide the ability to run concurrent dataplanes in parallel in order to support the dual dataplane model of CI/CD upgrades with inline validation.

(134) As illustrated in FIG. 5, ingress packets are received through port 0 **510** and port 1 **512** to the packet dispatcher **502**, and then the packet dispatcher **502** selects which of these data packets to direct to the primary dataplane **504** and the shadow dataplane **506**. For example, in the verification mode **406** the packets can be mirrored such that identical replicas are sent to both dataplanes. In contrast, in the scale-out mode **412**, different sets of data packets can be directed to each of the dataplanes. In FIG. 5 only the packets from the primary dataplane **504** are transmitted to port 0 **510** and port 1 **512**, which is reflective of the normal mode **404** and the verification mode **406**. In the scale-out mode **412**, however, both dataplanes will process and transmit packets from the network component/device implemented using the dual dataplane architecture **500**. Additionally, throughout the promotion mode **410**, the “New Primary” dataplane (i.e., shadow dataplane **506**) can gracefully transition from not sending packets to the ports to being the only of the two dataplanes that is sending packets to the ports. Control signal flow between the control-plane agent **508** and each of

the packet dispatcher **502**, the primary dataplane **504**, and the shadow dataplane **506**. Further, control signal can flow from the primary dataplane **504** to the shadow dataplane **506**.

(135) The dual dataplane architecture **500** includes a shared memory **516** that can be accessed by respective components of the dual dataplane architecture **500**, including, e.g., being accessed by the packet dispatcher **502**, primary dataplane **504**, shadow dataplane **506**, and control-plane agent **508**. The shared memory **516** enables the dataplanes to operate as stateless as possible by storing state-type information in the shared memory **516**, which is accessible by all components of the system. This can allow the packet dispatcher **502** to monitor the performance of both dataplanes. Further, these features can also allow flows to be migrated from one dataplane to the other since the state is isolated.

(136) The dual dataplane architecture **500** uses several functions to realize the various modes. For example, the network devices can be deployed on high availability (HA), and, more particularly, two modes of high availability (HA) can be used: (i) active/standby HA between the dataplanes (e.g., when in the Normal Mode active/standby HA is used with the primary dataplane in active HA and the shadow dataplane in standby HA) and (ii) active/active HA between the dataplanes (e.g., when in the Scale-out Mode active/active HA is used with both the primary dataplane and the shadow dataplane in active HA).

(137) For example, when the network device is a firewall, operating in active/standby HA entails that the first firewall processes all the traffic, and the second firewall, which is a clone of the first firewall, is waiting to take over. Continuing the non-limiting firewall example, operating in active/active HA entails both the first firewall and the second firewall are active (i.e., processing traffic). For example, half the traffic can be sent to the first firewall and the remaining half the traffic can be sent to the second firewall, thereby leveraging all of the compute on that system.

(138) Another function enabled by the dual dataplane architecture **500** is mirroring the packets to the shadow dataplane by sending replicas of the packets through both dataplanes. Because replicas of the packets are used through both dataplanes, this function is used in the verification mode to allow an apples-to-apples type comparison for verification testing between the current version of the network-device software, which is operating on the shadow dataplane. Based on this comparison, the system can tell if the new network software (or new network policy) satisfies predefined verification criteria and passes the verification test to replace the old network software (or old network policy).

(139) Another function enabled by the dual dataplane architecture **500** is gracefully transitioning flows from the “Old Primary” to “New Primary” dataplane (e.g., in the promotion mode).

(140) These functions are enabled (in part) by: (i) keeping the dataplanes as stateless as possible and (ii) storing state-type information in a shared memory that is accessible by all components of the dual dataplane architecture **500**. This allows the packet dispatcher **502** to monitor the performance of both dataplanes. This also allows flows to be migrated from one dataplane to the other because the state is isolated.

(141) Additionally, the dataplanes can share additional components. For example, the dataplanes can share a common random number generator. Consider the case in which the dataplanes are used when performing man-in-the-middle interception of TLS sessions. In this case, a secret key is generated, and the key should be the same value for both dataplanes so that the traffic can be mirrored through both data planes. That is, the dataplanes need to generate the identical key, so that they can understand the resulting response in the TLS session. The keys are generated using a random number from a random number generator. Accordingly, by using the same random number from a common random generator, the dataplanes can generate the same unique key for TLS or similar processes.

(142) FIG. **6** illustrates a non-limiting example of a DPU dual dataplane architecture **600**. To realize the various modes (e.g., normal mode **404**, verification mode **406**, promotion mode **410**, and scale-out mode **412**), the DPU dual dataplane architecture **600** uses the same (or similar)

functions as the dual dataplane architecture **500** to realize the various modes.

(143) For example, the network devices can be deployed on high availability (HA), and, more particularly, two modes of high availability (HA) can be used: (i) active/standby HA between the dataplanes (e.g., when in the Normal Mode active/standby HA is used with the primary dataplane in active HA and the shadow dataplane in standby HA) and (ii) active/active HA between the dataplanes (e.g., when in the Scale-out Mode active/active HA is used with both the primary dataplane and the shadow dataplane in active HA).

(144) Descriptions of the functions and how they enable the respective modes of the DPU dual dataplane architecture have been provided above in reference to FIG. 5. Therefore these descriptions are not repeated here. Some differences between the dual dataplane architecture **500** and the DPU dual dataplane architecture **600** are noted here. Generally, in the promotion mode, the graceful transition of packet flows from the “Old Primary” dataplane to the “New Primary” dataplane is enabled by keeping the dataplanes as stateless as possible and the storing state is provided in a shared memory that is accessible by all components of the system, thereby allowing the dispatcher to monitor the performance of both dataplanes. Further, these features can also allow flows to be migrated from one dataplane to the other since the state is isolated. In DPU environments, however, it may not be possible for the dataplanes to run entirely stateless, e.g., when leveraging hardware acceleration. To address this difference, during the promotion mode, session re-establishment may be used, resulting in a truly graceful promotion mode taking longer in the DPU dual dataplane architecture **600** than in the dual dataplane architecture **500**.

(145) The DPU dual dataplane architecture **600** includes a shared memory **620** that can be accessed by respective components of the DPU dual dataplane architecture **600**, including, e.g., being accessed by the packet dispatcher **602**, primary dataplane **604**, shadow dataplane **606**, and control-plane agent **608**. The shared memory **620** enables the dataplanes to operate as stateless as possible by storing state-type information in the shared memory **620**, which is accessible by all components of the system. This can allow the packet dispatcher **602** to monitor the performance of both dataplanes. Further, these features can also allow flows to be migrated from one dataplane to the other since the state is isolated.

(146) As illustrated in FIG. 6, ingress packets are received through port 0 **610a** and port 1 **612a** of DPU 0 **614**, and then the ingress packets are transmitted from the respective ports to the packet dispatcher **602**. The packet dispatcher **602** then selects which of these data packets to direct to the primary dataplane **604** and which to the shadow dataplane **606**. For example, in the verification mode **406** the packets can be mirrored such that identical replicas are sent to both dataplanes. In contrast, in the scale-out mode **412**, different sets of data packets can be directed to each of the dataplanes. In FIG. 6 only the packets from the primary dataplane **604** are transmitted to port 0 **610a** and port 1 **612a** of DPU 0 **614**, and the packets from the shadow dataplane **606** are transmitted to port 0 **610b** and port 1 **612b** of DPU 1 **618**.

(147) Depending on which mode the DPU dual dataplane architecture **600** is in, packets might not be sent to DPU 1 **618**, and, even when the packets are sent to and processed by DPU 1 **618**, the packets might not be transmitted from the ports of DPU 1 **618** (i.e., port 0 **610b** and port 1 **612b**). For example, in the normal mode **404**, the shadow dataplane **606** is in standby and does not receive any data packets. In contrast, in scale-out mode **412**, the shadow dataplane **606** is active, and both DPUs are engaged in processing and sending packets from their respective ports. Further, in the verification mode **406**, however, packets that are identical replicas are sent to and processed by both dataplanes, but only DPU 1 **618** transmits the processed packets from its ports.

(148) In the promotion mode **410**, there can be a gradual transition over time that starts with DPU 0 **614** transmitting the processed packets and end with DPU 1 **618** transmitting the processed packets. Accordingly, throughout the promotion mode **410**, the “New Primary” dataplane (i.e., shadow dataplane **606**) can gracefully transition from not sending packets to the ports to being the only of the two dataplanes that is sending packets to the ports.

(149) Control signals can flow between the control-plane agent **608** and each of the packet dispatchers **602**, the primary dataplane **604**, and the shadow dataplane **606**.

(150) FIG. 7A and FIG. 7B illustrate a non-limiting scaling up the system to more than two dataplanes. Generally, the total number of dataplanes can be any number N. In this example, the total number of dataplanes is N=8, but a person of ordinary skill in the art will recognize that any number of dataplanes can be used. When pairwise dataplanes are used, the number of dataplanes can be even, such that when in the verification mode **406** each primary dataplane is mirrored by a corresponding shadow dataplane. However, a primary dataplane can also be mirrored by more than one shadow dataplane (e.g., a first shadow dataplane that is implementing a first new version and a second shadow dataplane that is implementing a second new version).

(151) FIG. 7A illustrates pairwise dataplanes. when in a first configuration **700a** (e.g., a normal mode **404**, a verification mode **406**, or a promotion mode **410**), each primary dataplane is mirrored by a corresponding shadow dataplane. That is, when in first configuration **700a** primary dataplane 1 **702a**, primary dataplane 2 **702b**, primary dataplane 3 **702c**, and primary dataplane 4 **702d** are mirrored by shadow dataplane 1 **704a**, shadow dataplane 2 **704b**, shadow dataplane 3 **704c**, and shadow dataplane 4 **704d**, respectively.

(152) FIG. 7B illustrates a non-limiting example of a scale-out-mode configuration **700b** for eight dataplanes. When in the scale-out-mode configuration **700b** the primary dataplane 1 **702a**, primary dataplane 2 **702b**, primary dataplane 3 **702c**, and primary dataplane 4 **702d** can remain the same as in the first configuration **700a**. The shadow dataplanes are, however, converted to primary dataplanes, as in the dual dataplane examples discussed above. That is, shadow dataplane 1 **704a**, shadow dataplane 2 **704b**, shadow dataplane 3 **704c**, and shadow dataplane 4 **704d**, respectively become reassigned to primary dataplane 5 **706a**, primary dataplane 6 **706b**, primary dataplane 7 **706c**, and primary dataplane 8 **706d**, respectively.

(153) Thus, this system can be extended to any number N dataplanes where $(N \% 2 == 0)$. For example, in an implementation using a cluster of 8 DPUs is deployed in a top-of-rack switch, the system operates by creating four pairs of dual dataplanes.

(154) FIG. 8A illustrates an example of training an AI engine **808**. In step **804**, training data **802** (i.e., a corpus of verification metadata) is applied to train the AI engine **808**. For example, the AI engine **808** can be an artificial neural network (ANN) that is trained via supervised or unsupervised learning using a backpropagation technique to train the weighting parameters between nodes within respective layers of the ANN.

(155) In supervised learning, the training data **802** is applied as an input to the AI engine **808**, and an error/loss function is generated by comparing the output from the AI engine **808** with the desired output (e.g., a known prediction or label associated with the training data **802**). The coefficients of the AI engine **808** are iteratively updated to reduce an error/loss function. The value of the error/loss function decreases as outputs from the AI engine **808** increasingly approximate the desired output. In other words, ANN infers the mapping implied by the training data, and the error/loss function produces an error value related to the mismatch between the desired output and the outputs from the AI engine **808** that are produced as a result of applying the training data **802** to the AI engine **808**.

(156) Alternatively, for unsupervised learning or semi-supervised learning, training data **802** is applied to train the AI engine **808**. For example, the AI engine **808** can be an artificial neural network (ANN) that is trained via unsupervised or self-supervised learning using a backpropagation technique to train the weighting parameters between nodes within respective layers of the ANN.

(157) In unsupervised learning, the training data **802** is applied as an input to the AI engine **808**, and an error/loss function is generated by comparing the predictions of the next word in a text from the AI engine **808** with the actual word in the text. The coefficients of the AI engine **808** can be iteratively updated to reduce an error/loss function. The value of the error/loss function decreases

as outputs from the AI engine **808** increasingly approximate the training data **802**.

(158) For example, in certain implementations, the cost function can use the mean-squared error to minimize the average squared error. In the case of a of multilayer perceptrons (MLP) neural network, the backpropagation algorithm can be used for training the network by minimizing the mean-squared-error-based cost function using a gradient descent method.

(159) Training a neural network model essentially means selecting one model from the set of allowed models (or, in a Bayesian framework, determining a distribution over the set of allowed models) that minimizes the cost criterion (i.e., the error value calculated using the error/loss function). Generally, the ANN can be trained using any of numerous algorithms for training neural network models (e.g., by applying optimization theory and statistical estimation).

(160) For example, the optimization method used in training artificial neural networks can use some form of gradient descent, using backpropagation to compute the actual gradients. This is done by taking the derivative of the cost function with respect to the network parameters and then changing those parameters in a gradient-related direction. The backpropagation training algorithm can be: a steepest descent method (e.g., with variable learning rate, with variable learning rate and momentum, and resilient backpropagation), a quasi-Newton method (e.g., Broyden-Fletcher-Goldfarb-Shannon, one step secant, and Levenberg-Marquardt), or a conjugate gradient method (e.g., Fletcher-Reeves update, Polak-Ribière update, Powell-Beale restart, and scaled conjugate gradient). Additionally, evolutionary methods, such as gene expression programming, simulated annealing, expectation-maximization, non-parametric methods and particle swarm optimization, can also be used for training the AI engine **808**.

(161) The training **804** of the AI engine **808** can also include various techniques to prevent overfitting to the training data **802** and for validating the trained AI engine **808**. For example, bootstrapping and random sampling of the training data **802** can be used during training.

(162) Further, other machine learning (ML) algorithms can be used for the AI engine **808**, and the AI engine **808** is not limited to being an ANN. For example, there are many machine-learning models, and the AI engine **808** can be based on machine-learning systems that include generative adversarial networks (GANs) that are trained, for example, using pairs of network measurements and their corresponding optimized configurations.

(163) As understood by those of skill in the art, machine-learning based classification techniques can vary depending on the desired implementation. For example, machine-learning classification schemes can utilize one or more of the following, alone or in combination: hidden Markov models, recurrent neural networks (RNNs), convolutional neural networks (CNNs); Deep Learning networks, Bayesian symbolic methods, general adversarial networks (GANs), support vector machines, image registration methods, and/or applicable rule-based systems. Where regression algorithms are used, they can include but are not limited to: a Stochastic Gradient Descent Regressors, and/or Passive Aggressive Regressors, etc.

(164) Machine learning classification models can also be based on clustering algorithms (e.g., a Mini-batch K-means clustering algorithm), a recommendation algorithm (e.g., a Miniwise Hashing algorithm, or Euclidean Locality-Sensitive Hashing (LSH) algorithm), and/or an anomaly detection algorithm, such as a Local outlier factor. Additionally, machine-learning models can employ a dimensionality reduction approach, such as, one or more of: a Mini-batch Dictionary Learning algorithm, an Incremental Principal Component Analysis (PCA) algorithm, a Latent Dirichlet Allocation algorithm, and/or a Mini-batch K-means algorithm, etc.

(165) FIG. **8B** illustrates an example of using the trained AI engine **808**. The verification metadata **806** is applied to the trained AI engine **808** to generate various outputs, which can include the verification criteria **810**.

(166) For clarity of explanation, in some instances the present technology may be presented as including individual functional blocks including functional blocks comprising devices, device components, steps or routines in a method embodied in software, or combinations of hardware and

software.

(167) Any of the steps, operations, functions, or processes described herein may be performed or implemented by a combination of hardware and software services or services, alone or in combination with other devices. In some embodiments, a service can be software that resides in the memory of a client device and/or one or more servers of a system and performs one or more functions when a processor executes the software associated with the service. In some embodiments, a service is a program or a collection of programs that carry out a specific function. In some embodiments, a service can be considered a server. The memory can be a non-transitory computer-readable medium.

(168) In some embodiments the computer-readable storage devices, mediums, and memories can include a cable or wireless signal containing a bit stream and the like. However, when mentioned, non-transitory computer-readable storage media expressly exclude media such as energy, carrier signals, electromagnetic waves, and signals per se.

(169) Methods according to the above-described examples can be implemented using computer-executable instructions that are stored or otherwise available from computer readable media. Such instructions can comprise, for example, instructions and data which cause or otherwise configure a general purpose computer, special purpose computer, or special purpose processing device to perform a certain function or group of functions. Portions of computer resources used can be accessible over a network. The computer executable instructions may be, for example, binaries, intermediate format instructions such as assembly language, firmware, or source code. Examples of computer-readable media that may be used to store instructions, information used, and/or information created during methods according to described examples include magnetic or optical disks, solid state memory devices, flash memory, USB devices provided with non-volatile memory, networked storage devices, and so on.

(170) Devices implementing methods according to these disclosures can comprise hardware, firmware and/or software, and can take any of a variety of form factors. Typical examples of such form factors include servers, laptops, smart phones, small form factor personal computers, personal digital assistants, and so on. Functionality described herein also can be embodied in peripherals or add-in cards. Such functionality can also be implemented on a circuit board among different chips or different processes executing in a single device, by way of further example.

(171) The instructions, media for conveying such instructions, computing resources for executing them, and other structures for supporting such computing resources are means for providing the functions described in these disclosures.

(172) Although a variety of examples and other information was used to explain aspects within the scope of the appended claims, no limitation of the claims should be implied based on particular features or arrangements in such examples, as one of ordinary skill would be able to use these examples to derive a wide variety of implementations. Further and although some subject matter may have been described in language specific to examples of structural features and/or method steps, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to these described features or acts. For example, such functionality can be distributed differently or performed in components other than those identified herein. Rather, the described features and steps are disclosed as examples of components of systems and methods within the scope of the appended claims.

Claims

1. An apparatus for implementing a network component and verifying an update to the network component, the apparatus comprising: one or more ports configured to receive ingress traffic and transmit egress traffic; and circuitry comprising a first dataplane, a second dataplane, and a controller, wherein the first dataplane is configured to process, in accordance with first networking

instructions, a first subset of data packets from the ingress traffic to generate first egress data packets and a first record representing a performance of the first dataplane when processing the first subset of data packets, the second dataplane is configured to process, in accordance with second networking instructions, a second subset of data packets from the ingress traffic to generate second egress data packets and a second record representing a performance of the second dataplane when processing the second subset of data packets, the second subset being mirrored to the first subset such that the second subset is identical to the first subset, the controller is configured to receive and compare the first record with the second record and uses the comparison of the first record with the second record to verify an updated version of the network component, and the first networking instructions execute a current version of the network component and the second networking instructions execute the updated version of the network component.

2. The apparatus of claim 1, wherein the controller is configured to verify the updated version of the network component by determining that the comparison of the first record with the second record passes a verification test based on the comparison of the first record with the second record satisfying one or more predefined criteria.

3. The apparatus of claim 2, wherein the one or more predefined criteria of the verification test include that, for predefined parameters, values of the predefined parameters in the first record are within respective predefined ranges for the values of the predefined parameters in the second record.

4. The apparatus of claim 3, wherein the predefined parameters are selected from the group consisting of: (i) a minimum central processor unit (CPU) usage, (ii) a maximum CPU usage, (iii) an average CPU usage, (iv) a minimum memory usage, (v) a maximum memory usage, (vi) an average memory usage, (vii) a memory growth over a verification period, and (viii) a packet latency for packets to traverse a dataplane.

5. The apparatus of claim 2, wherein the circuitry is configured to determine the predefined criteria based on verification metadata, the verification metadata providing indicia regarding predicted differences between the performance of the first dataplane with respect to the performance of the second dataplane.

6. The apparatus of claim 5, wherein the verification metadata comprises learned values that are based on pre-verification testing the updated version of the network component in a test bed that simulates a commercial network, wherein the pre-verification testing occurs prior to the verification testing.

7. The apparatus of claim 2, wherein the network component implements a network policy, and the updated version of the network component represents a modification of the network policy relative to the current version of the network component.

8. The apparatus of claim 7, wherein the one or more predefined criteria of the verification test include: a first criterion that the updated version of the network policy does not adversely affect network traffic more than a predefined amount, and/or a second criterion that the updated version of the network policy provides an expected change to a predefined aspect of the network traffic.

9. The apparatus of claim 1, further comprising: a packet dispatcher that is configured to: receive the ingress traffic from the one or more ports, determine which of the data packets from the ingress traffic to include in the first subset and then transmit the first subset to the first dataplane, and determine which of the data packets from the ingress traffic to include in the second subset and then transmit the second subset to the second dataplane, wherein the circuitry comprises a first data processing unit (DPU) that implements the first dataplane and a second DPU that implements the second dataplane, and the packet dispatcher is implemented in only one of the first DPU and the second DPU, and the packet dispatcher spans data packets to the DPU of the first DPU and the second DPU in which the packet dispatcher is not implemented.

10. The apparatus of claim 1, wherein the first dataplane is configured to operate as a primary dataplane and the second dataplane is configured to operate as a shadow dataplane, such that the

primary dataplane provides a functionality of the network component and the shadow dataplane is only used to verify the updated version of the network component, and the egress traffic transmitted by the one or more ports comprises the first egress traffic and excludes the second egress traffic.

11. The apparatus of claim 1, wherein the circuitry comprises: one or more data processing units (DPUs) configured to use hardware acceleration to perform the first networking instructions and second networking instructions, or a virtual machine in a central processing unit (CPU) that is implemented as a Berkeley packet filter (BPF).

12. The apparatus of claim 1, wherein: the network component is configured to provide data-packet filtering, load balancing, security screening, malware detection, firewall protection, data-packet routing, data-packet switching, data-packet forwarding, computing header checksums, or implementing network policies.

13. The apparatus of claim 1, wherein the apparatus is an edge-computing processor that is configured in an embedded device of a network edge.

14. A method of implementing a network component, the method comprising: receiving, at one or more ports of a network device, ingress traffic comprising data packets; processing a first subset of data packets from the ingress traffic at a first dataplane processes, the first subset being processed in accordance with first networking instructions to generate first egress data packets and a first record that represents a performance of the first dataplane when processing the first subset of data packets; processing a second subset of data packets from the ingress traffic at a second dataplane processes, the second subset being processed in accordance with second networking instructions to generate second egress data packets and a second record that represents a performance of the second dataplane when processing the second subset of data packets, the second subset being mirrored to the first subset such that the second subset is identical to the first subset; comparing, at a controller, the first record with the second record and using the comparison of the first record with the second record to verify an updated version of the network component, wherein the first networking instructions execute a current version of the network component and the second networking instructions execute the updated version of the network component.

15. The method of claim 14, wherein the controller verifies the updated version of the network component by determining that the comparison of the first record with the second record passes a verification test based on the comparison of the first record with the second record satisfying one or more predefined criteria.

16. The method of claim 15, wherein the one or more predefined criteria of the verification test include that, for predefined parameters, values of the predefined parameters in the first record are within respective predefined ranges for the values of the predefined parameters in the second record, and the predefined parameters are selected from the group consisting of: (i) a minimum central processor unit (CPU) usage, (ii) a maximum CPU usage, (iii) an average CPU usage, (iv) a minimum memory usage, (v) a maximum memory usage, (vi) an average memory usage, (vii) a memory growth over a verification period, and (viii) a packet latency for packets to traverse a dataplane.

17. The method of claim 15, further comprising: determining the predefined criteria based on verification metadata, the verification metadata providing indicia regarding predicted differences between the performance of the first dataplane with respect to the performance of the second dataplane, and the verification metadata comprises learned values that are based on pre-verification testing the updated version of the network component in a test bed that simulates a commercial network, wherein the pre-verification testing occurs prior to the verification testing.

18. The method of claim 15, wherein the network component implements a network policy, and the updated version of the network component represents a modification of the network policy relative to the current version of the network component.

19. The method of claim 18, wherein the one or more predefined criteria of the verification test

include: a first criterion that the updated version of the network policy does not adversely affect network traffic more than a predefined amount, and/or a second criterion that the updated version of the network policy provides an expected change to a predefined aspect of the network traffic.

20. The method of claim 14, wherein the first dataplane is configured to operate as a primary dataplane and the second dataplane is configured to operate as a shadow dataplane, such that the primary dataplane provides a functionality of the network component and the shadow dataplane is only used to verify the updated version of the network component.
