

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12407654
Kind Code	B2
Date of Patent	September 02, 2025
Inventor(s)	Wang; Hongliang et al.

System and method for firewall policy rule management

Abstract

A system and method for updating network policies may include determining, from a set of hosts, a subset of hosts to have network policies updated; for each host in the subset of hosts, determining a set of policies relevant to the host, and determining a difference between the set of policies relevant to the host and a set of policies relevant to the host determined at an earlier time; and for each of the subset of hosts, installing the set of policies relevant to the host. The difference may be calculated or determined by comparing a file containing the set of policies relevant to the host, and a file containing the set of policies relevant to the host determined at an earlier time, e.g. determined the prior calendar day.

Inventors: Wang; Hongliang (Singapore, SG), Zhou; Rongfeng (Singapore, SG), Suzuki; Jack (Tokyo, JP), Wu; Jun (Shanghai, CN)

Applicant: Morgan Stanley Services Group Inc. (New York, NY)

Family ID: 94125730

Assignee: Morgan Stanley Services Group Inc. (New York, NY)

Appl. No.: 18/776556

Filed: July 18, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20250007884 A1	Jan. 02, 2025

Related U.S. Application Data

continuation parent-doc US 18344274 20230629 ABANDONED child-doc US 18461417
continuation-in-part parent-doc US 18461417 20230905 US 12224986 child-doc US 18776556

Publication Classification

Int. Cl.: H04L9/40 (20220101)

U.S. Cl.:

CPC H04L63/0263 (20130101); H04L63/0236 (20130101);

Field of Classification Search

CPC: H04L (63/0263); H04L (63/0236)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6460141	12/2001	Olden	N/A	N/A
6834301	12/2003	Hanchett	N/A	N/A
7676836	12/2009	Prigent	726/13	H04L 63/0227
7991899	12/2010	Schmidt	N/A	N/A
8904511	12/2013	O'Neill	N/A	N/A
10157290	12/2017	Sinha	N/A	N/A
11570148	12/2022	Liu et al.	N/A	N/A
2003/0021283	12/2002	See et al.	N/A	N/A
2003/0177389	12/2002	Albert	N/A	N/A
2005/0177829	12/2004	Vishwanath	N/A	N/A
2007/0271592	12/2006	Noda	726/1	H04L 9/3234
2008/0104661	12/2007	Levin	N/A	N/A
2009/0288135	12/2008	Chang et al.	N/A	N/A
2013/0227109	12/2012	Crudele	N/A	N/A
2014/0297774	12/2013	Munupalle	709/208	H04L 41/0889
2015/0324589	12/2014	Thanos	N/A	N/A
2018/0027006	12/2017	Zimmermann et al.	N/A	N/A
2018/0343236	12/2017	Pillay-Esnault	N/A	N/A
2020/0177548	12/2019	Devarajan et al.	N/A	N/A
2020/0296007	12/2019	Finn, II et al.	N/A	N/A
2020/0336513	12/2019	Martin	N/A	N/A
2020/0374268	12/2019	Wyatt et al.	N/A	N/A
2021/0136039	12/2020	Jin et al.	N/A	N/A
2021/0273977	12/2020	Karasaridis et al.	N/A	N/A
2022/0078619	12/2021	Mahamkali	N/A	N/A
2022/0109701	12/2021	Zeng et al.	N/A	N/A
2022/0166756	12/2021	Gupta	N/A	H04L 41/16
2022/0191248	12/2021	Pieczul et al.	N/A	N/A
2022/0286379	12/2021	Hofer et al.	N/A	N/A
2023/0042610	12/2022	Barton et al.	N/A	N/A
2023/0099259	12/2022	Kulaga et al.	N/A	N/A
2023/0156812	12/2022	Yang et al.	N/A	N/A

OTHER PUBLICATIONS

Notice of Allowance for U.S. Appl. No. 18/461,417 dated Nov. 21, 2024. cited by applicant

Primary Examiner: Little; Vance M

Attorney, Agent or Firm: PEARL COHEN ZEDEK LATZER BARATZ LLP

Background/Summary

RELATED APPLICATION DATA (1) The present invention is a continuation-in-part of prior U.S. application Ser. No. 18/461,417, filed Sep. 5, 2023, entitled “SYSTEM AND METHOD FOR FIREWALL POLICY RULE MANAGEMENT”, which in turn is a continuation of prior U.S. application Ser. No. 18/344,274, filed Jun. 29, 2023, entitled “SYSTEM AND METHOD FOR FIREWALL POLICY RULE MANAGEMENT”, each of which is incorporated herein by reference in its entirety.

FIELD OF THE INVENTION

(1) The present invention relates generally to firewall policy and rule management; for example to a scalable firewall policy and rule management solution.

BACKGROUND OF THE INVENTION

(2) Computer systems and technologies exist to provide network segregation, e.g., in an enterprise according to security, business or other requirements. In an environment or enterprise with tens or hundreds of thousands of computer hosts, such a network segregation requirement could result in millions of different network firewall policies or rules to manage. Hosts may be, e.g. end user computers (e.g. a desktop or laptop computer used by an enterprise employee), or servers, firewalls, or other computers operated by the enterprise. Each of these hosts may use hundreds or thousands of rules or policies, each rule or policy allowing, on a packet basis, connection (e.g. computer data exchange, communication or connection) to one other host: the default for hosts in such systems may be no connection is allowed to another specific host unless a specific rule allows for it. Each host may include an individually designed combination of rules or policies, each allowing for a connection from the host to another host; a single host may need to install a huge amount of network rules or policies. An installation or update of firewall rules or policies may cause problems in the affected hosts or network, and it may be difficult to diagnose problems caused by the installation or update of rules or policies, or identify which portions of the update or which rules cause the problem.

(3) Computer hosts may include virtual machines (VMs), and in some organizations hundreds or even thousands of VMs may be added or created and decommissioned every day. This may mean that the list of Linux or other hosts that need firewall rules is not static, but is dynamic. However, certain assumptions may hold, for example: 1) if a host needs firewall policy today, then it likely still needs a firewall policy tomorrow, e.g., the host is either newly created or still in use (although a firewall policy for a host may change day-to-day, or overnight); and 2) if a host does not need firewall policy today, then it likely will not need it tomorrow e.g., the host is decommissioned.

SUMMARY OF THE INVENTION

(4) A system and method may update network policies by determining, among a set of hosts, a subset of hosts to have network policies updated; for each host in the subset of hosts, determining a set of policies relevant to the host; and for each of the subset of hosts, installing the set of policies relevant to the host. The subset of hosts may be determined based on a category or division such as the time zone corresponding to the location of each host in the subset of hosts. The policies relevant

to the host may be received from a database and saved to a file with the set of policies relevant to the host; the host may then access the file.

(5) A system and method for updating network policies may include determining, from a set of hosts, a subset of hosts to have network policies updated; for each host in the subset of hosts, determining a set of policies relevant to the host, determining a difference between the set of policies relevant to the host and a set of policies relevant to the host determined at an earlier time; and for each of the subset of hosts, installing the set of policies relevant to the host. The difference may be calculated or determined by comparing a file containing the set of policies relevant to the host, and a file containing the set of policies relevant to the host determined at an earlier time, e.g. determined the prior calendar day.

(6) Maintaining a difference between a prior rule or firewall installation and a current installation (e.g. a difference in the state of a firewall between the period before and after an update) may allow the diagnosis of problems caused by the installation: for example, if the difference shows a set of rules (being a subset of all rules installed) were different between a certain day's update and a previous update, investigators may focus their diagnosis efforts on those rules that are different.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) Non-limiting examples of embodiments of the disclosure are described below with reference to figures listed below. The subject matter regarded as the invention is particularly pointed out and distinctly claimed in the concluding portion of the specification. The invention, however, both as to organization and method of operation, together with objects, features and advantages thereof, may best be understood by reference to the following detailed description when read with the accompanied drawings.

(2) FIG. 1 depicts a policy management system for a number of hosts in an enterprise, according to embodiments of the present invention.

(3) FIG. 2 shows a flowchart of a method according to embodiments of the present invention.

(4) FIG. 3 illustrates a series of operations of a calculate stage, according to some embodiments of the invention.

(5) FIG. 4 illustrates a series of operations of a simplify stage, according to some embodiments of the invention.

(6) FIG. 5 illustrates a series of operations of a difference stage, according to some embodiments of the invention.

(7) FIG. 6 illustrates a series of operations of an installer or query stage, according to some embodiments of the invention.

(8) FIG. 7 shows a high-level block diagram of an exemplary computing device according to some embodiments of the present invention.

(9) It will be appreciated that for simplicity and clarity of illustration, elements shown in the figures have not necessarily been drawn accurately or to scale. For example, the dimensions of some of the elements may be exaggerated relative to other elements for clarity, or several physical components may be included in one functional block or element. Reference numerals may be repeated among the figures to indicate corresponding or analogous elements.

DETAILED DESCRIPTION

(10) In the following detailed description, numerous specific details are set forth in order to provide a thorough understanding of the invention. However, it will be understood by those skilled in the art that the present invention may be practiced without these specific details. In other instances, well-known methods, procedures, and components, modules, units and/or circuits have not been described in detail so as not to obscure the invention. For the sake of clarity, discussion of same or

similar features or elements may not be repeated.

(11) Although embodiments of the invention are not limited in this regard, discussions utilizing terms such as, for example, “processing,” “computing,” “calculating,” “determining,” “establishing,” “analyzing,” “checking”, or the like, may refer to operation(s) and/or process(es) of a computer, a computing platform, a computing system, or other electronic computing device, that manipulates and/or transforms data represented as physical (e.g., electronic) quantities within the computer's registers and/or memories into other data similarly represented as physical quantities within the computer's registers and/or memories or other information non-transitory storage medium that may store instructions to perform operations and/or processes. Although embodiments of the invention are not limited in this regard, the terms “plurality” and “a plurality” as used herein may include, for example, “multiple” or “two or more”. The term set when used herein may include one or more items. Unless explicitly stated, the method embodiments described herein are not constrained to a particular order or sequence. Additionally, some of the described method embodiments or elements thereof can occur or be performed simultaneously, at the same point in time, or concurrently.

(12) Embodiments may divide, by time and by processing unit, the process of distributing network rules or policies from a central database to a large number of hosts (e.g. user computers, network devices, etc.) in an organization. One or more processes may determine, among the organization hosts, a subset to have network policies updated (e.g., a subset may be all hosts in a certain time zone, or another category of hosts), and for each of these selected hosts, may determine policies relevant or to be applied to each of the subset and store these policies, e.g. in a file, e.g. one file per host. Each host may then access its file and use the rules or policies in the file to update its network policies. In some embodiments, different processes or computers may perform the different tasks, improving such technology and resulting in scalability and the ability to perform this large task in parallel: e.g. a first process or device may determine the subset; one or a set of processes may access the policies for each host from a database and save to files; and another process may access the files on behalf of the host. In different embodiments the tasks need not be divided as in the examples provided herein.

(13) FIG. 1 depicts a policy management system for a number of hosts in an enterprise, according to embodiments of the present invention. Referring to FIG. 1, an enterprise **1** may include a number of hosts **10** (only one is shown for clarity), e.g. a server, firewall, user computer may include an installer agent or process **12** and a table **14**, e.g. an IPTable storing policies or rules **16**, e.g. Linux or other policies. Rules **16** may be used by firewall enforcement software to enforce a firewall relevant to the host on which they are installed, or to perform other tasks or enforcement. While one host **10** is shown for clarity, many hosts, e.g. many thousands, may be associated with or part of an organization or company. Installer process **12** may be an agent executing on host **10** with sufficient privilege (e.g., install privilege) to install or update policies or rules into table **14** when required (which may cause installation of policies to the associate host). An enterprise or other organization may include many, e.g. tens or hundreds of thousands, of hosts. CMDB (configuration management database) **20** may store information about hardware and software assets (commonly referred to as configuration items) including information describing hosts **10**. For example, CMDB **20** may store, for each of hosts **10** (e.g. the hosts in an enterprise), information such as hostname and location. Database **30** may store user-defined-network-policies, including rules and policies. Rules and policies in database **30** may be indexed to or associated with, in database **30**, hostnames, each hostname associated with a host **10**. One or more calculator services or processes **60** may access database **30** to determine a set of policies relevant to, to be applied to, or associated with a specific host, receive or extract policies or rules relevant to that host, create or maintain directories to store policies, and create a network policy file **42** (or other suitable data object) storing those policies. Calculator services **60** may be distributed geographically across enterprise infrastructure, but may be centrally located. Calculator services **60** typically work closely with database **30**; if

database **30** has a replica in a region then a calculator service **60** may be geographically distributed to that region. Each file **42** may be compressed. Orchestrator **70** may identify or determine which hosts, e.g. a subset of hosts from the set of hosts associated with an enterprise or organization, are to have network policies or rules updated at a specific time or time period (e.g., based on the time zone of the hosts, or another characteristic of the hosts). Shared file system **40** may include sufficient space to store files **42**.

(14) A cache **80** may perform services such as determining if policies for a certain host have been calculated in the past. Simplify services **82** may remove or strip away rule information not needed for certain functionality. Diff services **84** may compare policies or rules to determine a difference between a set of policies and a prior period set of policies. For example, diff services **84** may each compare a file containing the set of policies relevant to a specific host, and a file containing the set of policies relevant to that host determined at an earlier time.

(15) In one embodiment, shared file system **40** is organized using a directory structure where the names of the directories and files indicate the time period relevant to the stored policies (e.g. 2023 Dec. 31 being the example name of a directory storing policies created on Dec. 31, 2023) and/or the host for which the policies are relevant (e.g. hostname_1 being the example name of a directory storing policies relevant to a host named hostname_1). The directory storing rules or policies indexed for or relevant to a certain host may be indexed or described by that host, and may also be indexed or described by the time period (e.g. a cycle time period) relevant to the policies. A directory may include a set of files including, for example a file indicating policy differences between policies created in a most recent time period and those created in a previous time period. In another embodiment, such a directory structure need not be used. For example, each file **42** may be a temporary file created by a calculator process **60**, as described herein, storing policies extracted from database **30** and relevant to a specific host **10**.

(16) Shared file system **40** may be shared among entities such as calculator services **60** and reader service **52**, and other entities such as some system administrators. Reader service **52** and other entities such as system administrator may have read-only access to each file **42**.

(17) In some embodiments, a file **42** created for a host persists for a certain fixed amount of time (e.g. 24 hours), and then is deleted, and not reused for another host; however, different embodiments may use other methods of using such files. For example, a calculator service may perform calculations at 1:00 AM local time for a host, and an installer agent may install at 2:00 AM local time. In such a system, the host file may be refreshed every 24 hours. In other embodiments files **42** and other files such as difference files may persist for a longer period of time.

(18) A reader service or process **52** may be used by agent or process **12** to fetch or access the network policy file associated with the host executing agent **12**, e.g. from the shared file system. In FIG. **1** one host **10** is shown using reader service **52** to access one file **42**; in practice many hosts may each use a reader service to access a network policy file **42** concurrently.

(19) Hosts as used herein may be computers (e.g. desktops, servers, laptops, smartphones, etc.) and may be Linux hosts, e.g. operating the Linux system; however, systems other than Linux may be used.

(20) While data flows or connections in FIG. **1** are indicated by lines with arrows, in general the various entities in FIG. **1** such as hosts **10** and reader service or process **52** may communicate and exchange data via networks **2**, such as intranets (in some embodiments networks **2** do not include the Internet, as communications are typically internal to an organization; however in other embodiments the Internet may be used). Further, entities in FIG. **1** may exchange data on message bus **6**, for instance, the Apache Kafka event store and stream processing platform; message bus **6** may make use of networks **2**. For example, message bus **6** may be used to communicate identities of a subset of hosts to have network policies updated.

(21) One or more servers such as servers **4A** (e.g. hosting a shared file system), **4B** (e.g. hosting calculator, cache, simplify, diff and orchestrator processes) and **4C** (e.g. hosting databases and a

message bus) may host and execute processes such as shared file systems, calculators, orchestrators, databases, etc. While specific elements of FIG. 1 are hosted by specific servers, in other embodiments, such elements may be hosted by other servers, in cloud based systems, etc.

(22) Each host **10** may be configured such that communication with another host or computer is not permitted unless a specific rule of rules **16** permits communication, e.g. an exchange of packets, with a specific host or computer. Table 14 may be a table or other data structure (or multiple tables or data structure), accessed by a user-space utility program that allows a user or system administrator to configure rules **16** such as IP packet filter rules for the host, e.g. the rules of a Linux kernel firewall for the host, but may be another type of rule storage structure, and need not use the Linux environment. Rules **16** may be in one or more tables 14, and may tell the host's system how to treat network traffic packets. A host may include many, e.g. thousands of, rules **16**. Each of rules **16** may enable connection to another host, or provide another function, and may take the form of, for example—A OUTPUT-p tcp-d 1.2.3.4/23-j ACCEPT.

(23) Rules **16** may govern connections between computers, or user or other log in procedures, such that a mistaken or misguided rule installation may, in the case of high network traffic, prevent logins or cause disk full conditions. Rules **16** may govern connections between applications and hosts: for example, an app executing on host A connecting to computer B may correspond to a rule allowing for this, and a mistakenly installed rule or erroneous rule may prevent such a connection. Other rule formats may be used.

(24) Installer process **12** may be subject to restrictions such as it is only allowed to execute at certain times (e.g., at 2:00 AM local time for its associated host) so that it has least impact on user's daily work.

(25) A user **54** may use GUI (graphical user interface) **50** to access reader process **52** to review rules or policies assigned to a host. The user may enter the specified hostname into GUI **50** which may contact reader service or process **52** to fetch the corresponding file from the shared file system.

(26) The system depicted in FIG. 1 may include multiple processors, as each of computers and modules such as servers **4A**, **4B**, and **4C**, hosts **10**, installer agent or process **12**, calculator services or processes **60**, orchestrator **70**, and other modules may be, or include, components such as in FIG. 7. Thus, in practice, among a number of processors (typically each associated with a memory), a processor among the processors may determine a subset of hosts to update; another processor (possibly the same processor as the first in some embodiments) among the plurality of processors may determine a set of policies relevant to hosts, perform difference calculations, etc.; and another processor (possibly the same processor as the first and/or second in some embodiments) may install a set of policies relevant to a host. Other combinations of processors and functionalities may be used.

(27) Embodiments may improve the technologies of policy and firewall management by allowing for transparency and scalability of firewall management in large organizations. In a small organization with few hosts to manage, a server and database may provide policies for deployment to hosts. In large scale host management this does not work as well, as each of thousands of hosts may have thousands of policies individually specified for that host, meaning tens of millions of policies may need to be distributed to tens of thousands of hosts. Each policy may control or allow access between the host on which it is installed (in some cases the default is that absent a policy, the host cannot connect to another host), and possibly one other host; thus in such a system each host may need many policies. Embodiments may improve the distribution of such policies, and transparency, by precomputing or pre-filtering network policies relevant to a host, and providing the pre-filtered policies in a file or shared folder or file system, to be fetched by the destination host. By placing relevant policies in a file, and through the use of a shared file system, embodiments may provide transparency, allowing a user to easily see the relevant policies for auditing, troubleshooting, or other purposes. Prior art systems keep such policies in a database,

which is more difficult to access and less scalable. As the total number of hosts grows, the tools to update policies (e.g. calculator services or processes **60**) may grow with them, providing scalability. This is especially important when hosts are distributed geographically. Embodiments may solve scalability problems by splitting computational targets, e.g. of which hosts get which policies, and database access, from all-at-once across a global system of hosts, to smaller regional needs (e.g. dividing by time zone) and making use of asynchronous communication (e.g. such that policies are fetched from a database, stored in a file, then accessed by a host). In one embodiment, all firewall policy computations for a region may be completed within one hour. Embodiments may improve firewall or policy handling by allowing a history of what policy was installed when to be determined, e.g. via a diff file.

(28) Embodiments may improve the efficiency of computer management technology by not requiring applications to query a centralized database to retrieve Linux host information; rather embodiments may rely on past record to estimate host information, and allow easy troubleshooting of network connectivity issues, since often networking problems result from a network policy being wrongly placed onto a host, e.g. overnight during low use or non-use. In prior approaches host agents directly query a database to fetch network policies, which does not scale: with for example 100,000 hosts and 1,000,000 network policies, a database may become a bottleneck and may result in lagging in network policy retrieval. Embodiments of the present invention may improve on prior art technology with a scalable method by using asynchronous communication, and precomputing of network policies. In one embodiment, when a host needs to refresh its network policy, it may go to a shared disk or data repository to fetch the precomputed result instead of querying a database directly. This may eliminate a bottleneck, resulting in scalability. Embodiments may make network troubleshooting more efficient, as employees may easily determine policy changes, e.g. which may occur overnight.

(29) FIG. 2 shows a flowchart of a method according to embodiments of the present invention. The operations of FIG. 2 may be performed by the systems shown in FIG. 1 and other figures shown herein, but other systems may be used. The operations of FIG. 2 may be performed repeatedly, e.g. once per a cycle time period, e.g. once a day.

(30) Referring to FIG. 2, in operation **100**, a process may identify or determine a subset of hosts or computing devices, from a larger set of hosts (e.g. all hosts in an organization) to have network policies updated or distributed to, from a set of hosts. This may be termed a calculate process. For example, an orchestrator service such as orchestrator **70** may execute or be triggered, for example periodically, such as an hourly basis. A set of hosts corresponding to a sub-time period within a cycle time period, such as a one-hour time period during an “off peak” time period, e.g. 2-3 AM, may be selected, where the set of hosts resident in the sub-time period changes as the sub-time period occurs at different times around the world. When executed, an orchestrator service may, for example, fetch Linux or other host information from a data store such as CMDB **20**, and filter or select those hosts which fit a certain criteria, or receive an appropriately filtered list from a data store. For example, the subset of hosts may be determined based on the time zone corresponding to the location of each host in the subset of hosts: hosts that are, per the data store, located or physically present at a certain local time (e.g. local time 1:00 AM) may be selected. For example, all Linux hosts within region Singapore and Hong Kong may be selected as part of one subset, as these locations are in time zone GMT+8.

(31) The selected subset of hosts (e.g., the identities or identifiers identifying these hosts) may be then transmitted to other processes; for example an orchestrator service may publish the selected Linux hostnames onto message bus **6**. In such an example, where all hosts within an organization that reside within a certain time zone are updated at approximately the same time, e.g. at local time 1:00 AM, it may create a natural division by hour to reduce the computer infrastructure burden on accessing update data and updating hosts. While time zones are used as an example method of categorizing or dividing hosts into groups or subsets for updating, in other embodiments other

methods may be used. In one embodiment, an orchestrator service automatically starts execution periodically, e.g. every hour, and fetches all hosts identifications from a database such as CMDB **20**, then filters and selects hosts in a certain category (e.g. located where local time is 1:00 AM, based on host location information); then publishes the filtered or selected hostnames, e.g. onto a message bus one by one. An orchestrator process may include a list of locations correlated with locations, and use the time of its execution to determine which location to use to filter or select hosts; e.g. an orchestrator may correlate GMT+8 with Singapore and Hong Kong. A sample message published by an orchestrator is {'hostname': 'abc'}, where abc is an enterprise Linux hostname, and each message represents one host. Other methods of transmitting selected hosts to other processes may be used.

(32) In operation **110**, for each host in the subset of hosts or computing devices, a process may determine or identify a set of policies or rules relevant to or associated with that host. For example, a calculator service or process **60** may access and consume an individual hostname produced in operation **100** from a message bus, match or correlate the hostname to entries for network policies from a database storing network policies, e.g. by querying database **30** with a hostname, and receive from the database policies or rules relevant to the host. Multiple calculator services or processes **60**, each servicing one host at time, may exist and operate substantially concurrently, but these processes may be reused for multiple hosts, after each host is serviced. Calculator processes or services may operate by periodically consuming messages from a message bus, e.g. each calculator process consuming one message every second, process it (e.g. identify policies relevant to a hostname) and move onto the next message. Thus in one embodiment calculator services (and possibly other services discussed herein) are reused for different hosts.

(33) In one embodiment, calculation for a set of hosts occurs before (e.g. one hour before) queries for rules or policies are sent or are expected to be sent from hosts. Calculation may be based on a host list prepared in the prior cycle time period, e.g. the previous day. Other methods of determining policies relevant to a host may be used.

(34) Rules or policies relevant to or to be applied to a host may be those having in the rule or policy the hostname for that host; rules or policies that are designed to work with only that specific host (e.g. determining which other hosts the relevant host may connect with); being indexed to that host (e.g. being stored in a database with an identifier or index which correlates to the host or hostname); rules or policies that are specifically relevant to a single host by sharing a hostname with the host (e.g. in the form of a rule index or label); or those rules or policies otherwise being specified as being used for that particular host. In some embodiments, each rule or policy is indexed or mapped itself, individually, to one unique host and is not applicable to multiple different hosts; however, other indexing or relevancy schemes may be used. Such indexing or association with a host may be part of the rule itself.

(35) Each of a plurality of processes may operate substantially at the same time, at overlapping time periods, or concurrently to determine a set of policies relevant to the subset of hosts. Policies for multiple hosts may be determined concurrently, and multiple calculator or other services may operate concurrently. Thus, in one embodiment, all hosts located in one time zone, or one category, may have their policies determined at the same time or concurrently. Other methods of determining policies to apply to hosts may be used; e.g. a calculator service need not be used.

(36) In one embodiment, a calculator service may communicate policies or rules to a host via a calculator service output. A calculator service may produce, or save the policies to, a file or other data object (e.g. file **42**) with the set of policies relevant to the host, which may be a compressed file with computed network policies for a given host, where the file is stored on a shared file system, e.g. on a server such as server **4A** or a cloud service. Other methods of transmitting identified policies to hosts may be used.

(37) In operation **120**, a difference may be determined or calculated between the set of policies or rules relevant to or associated with the host and a previous set of policies or rules relevant to or

associated with that host, which were determined at an earlier time. For example in the case that the periodicity of determining relevant policies is once a day such that a cycle time period is 24 hours, the current or latest policies determined for a specific host in operation **100** may be compared against those determined for that host the previous day (e.g. the immediately preceding cycle time period). In some embodiments, the earlier time may be two time cycles in the past—e.g. two days ago. In some embodiments, a difference file may be saved in the same folder or location as the calculated policies. The difference may be determined by comparing a file containing the set of policies relevant to the host, and a file containing the set of policies relevant to the host determined at an earlier time, e.g. a previous file. The difference file may be saved in a directory or folder indexed by the host and time period, along with one or more files including the set of policies or rules relevant to the host.

(38) In operation **130**, further processing on calculated rules or policies may be performed. For example, a simplify process may remove or strip away rule information not needed for certain functionality.

(39) In operation **140**, for each of the subset of hosts, the set of policies relevant to or specific to the host may be installed in that host. For example, at a certain standardized time, e.g. 2:00 AM local time to the host, an installer agent or process **12** executing on a host may wake up or be executed and access a file or other data store, e.g. the most recently created file, of rules or policies for the host on which it is executing. E.g., a host, using such an installer process, may contact reader service or process **52** (e.g. executing on a server or cloud system separate from a host) to access or fetch its own network policy file (e.g. associated with or relevant to its host) from a shared file system or another file system, to access policies relevant to the host. It may decompress the file and install accordingly to update policies, for example by replacing, updating or modifying a table such as table 14. The installer process may use, e.g. an iptables-save command to install the rules; other methods may be used. This may occur at a regular time, e.g. in coordination with or shortly after operations **100** and **110**. For example, in one implementation, Singapore and Hong Kong servers (as hosts) start to fetch and install network policies at 2:00 AM local time. In some embodiments, a separate reader service need not be used to access a file of rules. In one embodiment, one reader service exists for an organization, servicing thousands of hosts, where the hosts contact the reader service by HTTP request.

(40) A reader service may provide security, authentication, authorization and auditing, and prevent impersonation and information leaks. For example, a host agent calling a reader service may transmit IP address, hostname and a secret key to the reader, which may then verify this information. If verification is successful, the reader returns the correct firewall to that host; if unsuccessful, the reader does not return anything to the host, and may log a warning message and/or raise an alert for enterprise security team to follow up. If a human system administrator calls the reader, the reader may verify the person's identity and permissions, and only if verified does the reader return information re hosts the person can access. For example, in some embodiments a system administrator in department A cannot view firewall rules of hosts in department B.

(41) In operation **150**, hosts may operate according to policies. For example, a rule or policy may, when processed by a program such as the iptables utility, prevent or allow, on a packet basis or another basis, connection (e.g. computer data exchange, communication or connection) to one other host. In some embodiments, the default for hosts in such systems may be no connection is allowed to another specific host unless a specific rule allows for it. Rules may have other or different effects.

(42) In operation **160** action may be taken based on the difference files, such as reversing or unwinding a rule or policy installation. For example, in case a user needs to reverse firewall rules to be those of a past date or cycle time period, they may specify that date or period as an input parameter for example to a Linux agent. The agent may in turn contact a reader service to fetch the ruleset from that period accordingly, e.g. from a directory with the appropriate date, and install the

rules or policies.

(43) Other operations or series of operations may be used.

(44) While specific modules and servers are described herein as executing specific functions, in other embodiments, the functions described herein may be performed by different modules or hardware.

(45) An embodiment may store, for example at filesystem **40**, database **30**, or on another system, firewall policy differences between different dates or updates, for example in difference files or another storage structure. In one embodiment, differences may be stored in a shared folder and/or filesystem. For example, in an embodiment where policies are updated daily, a filesystem folder labelled with the date of the update (e.g. /2023-12-30) may store all difference files for that date, and subdirectories for each host (e.g. /2023-12-30/hostname_1; /2023-12-30/hostname_2; etc.) may store, for each host corresponding to or relevant to the subdirectory, files recording the differences of policies between the date corresponding to the main directory and a prior update of policies. In the case that these directories accumulate increasing amounts of directories and files, directories and/or files may be periodically deleted. While in some embodiments, the date is used to calculate and label differences in rule or policy updates, other time markers (e.g. time and date; time; month; etc.) and markers other than time, may be used to calculate and label differences.

(46) A subdirectory for a host may include, for example: a file including all rules or policies for that host (without considering differences), which may be compressed, e.g. using ZIP or other compression; a file with a human-readable description of each rule (e.g. “allow a connection between host A and host B”); and one or more difference files showing the differences of policies between a date such as the date corresponding to the main directory and the date of a prior update of policies. A file showing the difference between rule or policy set A (e.g. in a file) and rule or policy set B (e.g. in a file) may indicate the rules or policies that are changed, new, or deleted in set B compared to set A. These are typically relevant to or applicable to a particular host: for example the differences may show a difference between the set of policies relevant to a certain host and a set of policies relevant to that same host determined at an earlier time. For example, one difference file may record only rules or policies that are new, deleted or updated over the last day (e.g. when comparing a current set of relevant policies with a set of relevant policies determined at an earlier time), and a second difference file may record only rules or policies that are new, deleted or updated over the last two days (in the case that an update over a weekend or short holiday causes problems that are not apparent until the next working day). Other information may be stored; ranges of days of difference beyond two may be used; and other formats may be used, with other directory structures, or not relying on directory structures.

(47) An embodiment may include a calculate stage, a query stage, a preparation stage and a clean stage.

(48) A calculate stage can be divided into steps such as: calculate, simplify and diff (e.g., compute a difference). FIG. 3 illustrates a series of operations of a calculate stage, according to some embodiments of the invention. Elements in FIG. 3 may be analogous to like-numbered elements in FIG. 1 and other figures shown herein. Referring to FIG. 3, an orchestrator service **70** may be executed or triggered on an hourly basis and may have a list of hosts that will be fetching firewall rules for next period of time, e.g. hour. Orchestrator service **70** may publish selected (e.g., selected as being in a certain time zone, or relevant to a particular sub-time period) Linux hostnames (e.g. one by one) onto message bus **6**. A sub-time period that determines a subset of hosts may be based on region. For example, all hosts in a certain region (regardless of time zone) may be updated, for example all hosts residing in New York and New Jersey, or a continent, such as Asia, Europe or North America. Thus hosts in the same region such as continent may have rules calculated at the same time even if they are in different time zones. In such an embodiment, policy or firewall fetch time may still be a certain specific time range per subset of hosts within the region, e.g. 2:00 AM per time zone, or another time having minimum impact on users' daily work.

(49) One or more calculator services **60** may consume individual hostnames from message bus **6**, and match network policies from database (for example, network policies matching hostnames may be extracted). The output of each calculator service **60** may be a firewall file, which may be compressed. Compression may be in the gzip format, or other formats. A two-level directory file structure may be created (e.g. by a calculator service) to store firewall files, and may be stored, for example, in shared file system **40**. The top-level directory may be named to indicate the current date, and the second level directory may be named to indicate the hostname. For example, if the hostname for a file (e.g. the file includes rules relevant to that hostname) is hostname_1 and the current date is Dec. 31, 2023, then the output directory and file may be/2023-12-31/hostname_1/firewall.gzip.

(50) An embodiment may operate on cycle time periods and sub-time periods, where a process repeats during each cycle time period, and each cycle time period is divided into sub-time periods, where different hosts have policies processed and updated during sub-time periods. For example, a cycle time period may be a day, a 24-hour period, which may be divided into sub-time periods of one hour. In a global system, where at any one time a set of hosts are resident in a time zone which has a local time being a certain update time (e.g. 2:00 AM), the cycle time period of a day may be divided into one-hour sub-time periods corresponding to local time. For example, a set of hosts may have policies updated during the sub-time period corresponding to 2:00-3:00 AM local time.

(51) FIG. 4 illustrates a series of operations of a simplify stage, according to some embodiments of the invention, which may remove or strip away rule information not needed for certain functionality. Elements in FIG. 4 may be analogous to like-numbered elements in FIG. 1 and other figures shown herein. Referring to FIG. 4, a generated firewall file (e.g. firewall.gzip) may contain both policy information and IP information. However, to facilitate production support troubleshooting all that may be needed is the policy information. Sample policy text is shown in Table 1 below: #[63935684] ec_core_infra_tcp_80 A PRODPERIM_OUTPUT-d 10.113.70.10/31-j ACCEPT

Table 1

(52) Multiple entries (e.g. millions in some cases) as in Table 1 may be included in a firewall.gzip file. In Table 1, the beginning “#” of the first line indicates this line is a comment. 63935684 is policy number, cc_core_infra_tcp_80 is policy or rule name, and ‘A PRODPERIM_OUTPUT-d 10.113.70.10/31-j ACCEPT’ is the actual IP table rule, which in this specific example is compatible with Linux command iptables. Systems attempting to diagnose which rules or rule changes caused a problem may be more efficient without extraneous information (which may be removed by a simplify process). For example, diagnosis may only require a policy number (e.g. 63935684) and a policy name (e.g. cc_core_infra_tcp_80). Thus a simplify service may input or consume the output of a calculator service (e.g. a firewall.gzip file), and remove or filter out all information such as in Table 1 except for policy or rule number and name. A simplify service may input information as in Table 1 and output, or collect in a file, only “#[63935684] ec_core_infra_top_80” from the input of Table 1 (but may do this for many entries such as in Table 1). A calculator service **60** may publish a firewall file path onto message bus **6** (e.g. with a message bus topic of “#2: simplify”) and this may be consumed by one or more simplify service processes **82**. More than one simplify service **82** instance or process may be used, e.g. concurrently, to service multiple hosts, each having a rule set to process. The output of a simplify service **82** may be a file (e.g. a text file, or another file) including the rule numbers and rule names for policies installed on the specific host corresponding to or named in the file path, without other information that might be included in a firewall.gzip or other file). Such a file may be placed or stored into the same directory as the firewall file (e.g. containing rules or policies) corresponding to the host corresponding to (or relevant to) the simplify output file, e.g. “policy.txt.” For example, if the relevant hostname is hostname_1 and the current date is Dec. 31, 2023, then the output file may have a path/filename of/2023-12-31/hostname_1/policy.txt. Such a file may be a text file. While text files, and files in .txt format, are

used as examples herein, files in other formats may be used.

(53) FIG. 5 illustrates a series of operations of a difference stage, according to some embodiments of the invention, which may for example determining a difference between a “current” or most recent set of policies relevant to a host and a set of policies relevant to the host determined at an earlier time, e.g. in a previous time period or day. Elements in FIG. 5 may be analogous to like-numbered elements in FIG. 1 and other figures shown herein. Such a stage may be needed to compare firewall policy differences between different installations (e.g. installed at different times), or between different dates or other times. For example, a user may have used their computer to access a host yesterday, but today she cannot use the computer to access that host. This may be due to firewall rules which have changed, for example overnight. A simplify service **82** may publish a policy or rules onto message bus **6** (e.g. using “topic #3: diff”) and this may be consumed by one or more difference services **84**. A difference service **84** may compare two sets of policies or rules (e.g. each set in a file, although other storage formats may be used), to calculate or determine a difference between the current or most recent set of policies relevant to the host and a set of policies relevant to the host determined at an earlier time. The difference may be determined by comparing a file containing the set of policies relevant to the host, and a file containing the set of policies relevant to the host determined at an earlier time. For example, the utility “diff”, a data comparison tool, may be used to compute differences between the contents of files.

(54) The output of a difference service **84** may be a file, e.g. a text file, including or indicating differences between a current policy (e.g. the latest policy or set of rules to be installed) and a previous policy (e.g. installed at previous times or dates), typically for one specific host only. This difference file may be placed or stored into the same directory as the firewall file (e.g. containing policies or rules relevant to the host). In one embodiment the difference information is saved in a first file, the (e.g. current) set of policies relevant to the host is saved in a second file, and the first file and the second file are stored in the same directory.

(55) The directory storing the current policies, the difference information, and possibly other information such as simplified policy information may be indexed by the host relevant to the policies (e.g. “hostname_1” may be included in a directory name) and the time period relevant to the policies (e.g. “2023 Dec. 31” representing Dec. 31, 2023, may be included in a directory name). An example file name for a difference file includes <date>_<previous_date>.diff. For example, in the case that a hostname is named hostname_1 and the current date is Dec. 31, 2023, then the output path and file, corresponding to the hostname, may be 2023 Dec. 31/hostname_1/2023-12-31_2023-12-30.diff. Other nomenclatures for allowing a path or filename to correspond to a host or hostname may be used.

(56) In one embodiment, multiple difference services **84** each consume files to compare from message bus **6** and produce for each comparison a diff file. Typically, the diff files are not distributed to the relevant hosts; rather they may be stored in a central directory structure, e.g. at server **4A**.

(57) Example contents of a diff or difference file are shown in Table 2 below: \$ cat 2023 Dec. 31-2023-12-30.diff 0a1 >25154576, Proxy Logging-Inbound

Table 2

(58) The contents of Table 2 indicate that over the night of Dec. 30, 2023, a policy with id=25154576 and name=“Proxy Logging-Inbound” was installed on this particular host, which is the host corresponding to or relevant to the file.

(59) FIG. 6 illustrates a series of operations of an installer or query stage, according to some embodiments of the invention. Elements in FIG. 6 may be analogous to like-numbered elements in FIG. 1 and other figures shown herein. Referring to FIG. 6, at operation **601**, at a certain time (e.g. the convenient time of 2:00 AM local time), installer agent **12** on a Linux host **10** may wake up and contact reader service **52** to fetch its own corresponding network policy file from the shared file system **40**.

(60) At operation **602**, upon receiving such request, reader service **52** may first consult a cache service **80** to check if a corresponding firewall or network policy has already been calculated or not. At operation **603**, if one has (yes), reader service **52** may combine data such as a current period of time (e.g. current date, and/or current time) and hostname to find the appropriate location (e.g. a directory indexed by time period and hostname) of the relevant firewall file and return the content to or installer agent or process **12** or Linux host **10**. In the meantime, reader service **52** may record a note or indication that this host is already queried, or already has made a query for a certain update cycle (e.g. for that day). In one example use case, a database table named, for example, query_record and has three columns: hostname, queried and is_freshman. “Freshman” may indicate a host that is newly seen to the various systems. Reader service **52** may execute a database command, such as the example SQL command such as: “UPDATE query_record SET queried=1 WHERE hostname=<host>”. Such a database may be global, recording information on all hosts in an organization, or a different database may be used for each sub-time period. Such a table or tables may be maintained, for example, in CMDB **20**.

(61) Embodiments of the invention may account for the case where, in-between the times when a set of hosts (e.g. hosts residing in a certain time zone) is to have rules updated, hosts are added to that set of hosts (e.g. newly installed in an organization) or deleted. Newly added hosts may not be recorded by reader service **52**, and reader service **52** may include un-needed entries for hosts that have been deleted.

(62) Reader service **52** may query cache service **80** to determine if a host has rules computed: if cache service returns yes for a certain host (that host has rules computed), then reader service **52** may perform SQL=UPDATE query_record SET queried=1 WHERE hostname=<host>.

(63) If cache service **80** returns no, then reader service **52** may execute SQL=INSERT INTO query_record (hostname, is_freshman) VALUES (‘<host>’, 1). If, during a cycle of rule updates (e.g. once per cycle time period), a host known to cache service **80** as being assigned to the currently being updated time zone or other grouping (e.g. during a and sub-time period) does not make a query, it can be assumed to be deleted, and may be deleted from cache service **80**. Thus, in one embodiment, the determining from a set of hosts, a subset of the hosts to have network policies updated may be based on a table (e.g. maintained at reader service **52**), where a host may be added to the table if a query is received from the host and the host is not already in the table; and a host may be removed from the table if a query is not received from the host during a predetermined time period or cycle time period (e.g. during the 24 hour period or cycle time period prior to a calculation being performed for the hosts in the relevant time zone).

(64) After operation **603**, if cache service **80** determines the firewall has not been calculated for this host, then reader service **52** may calculate a firewall (e.g. determine rules or policies for the relevant host), possibly on the fly, in real time (e.g. as opposed to being precomputed); and return a result to an installer agent or process **12**; and save a record of a hostname into database **30**, so that in the next cycle policies for that hostname can be precomputed. A record may be saved to database **30** because one embodiment works on the assumption that if a host needs firewall policy today, then it likely will exist and need that same policy tomorrow, or the next relevant time period. An example corresponding database command, such as an SQL command, may be: “INSERT INTO query_record (hostname, is_freshman) VALUES (‘<host>’, 1)”.

(65) At operation **604**, once a host **10** or installer agent or process **12** receives the result from reader service **52**, it may then decompress that file and install policies relevant to the host accordingly on host **10**. For example, in one example implementation, all Cal and Hong Kong resident hosts **10** (e.g. servers) may start to fetch and install network policy at 2:00 AM their local time.

(66) Operations **601-604** may repeat for a period of time (e.g. the next hour in the case a global system of hosts is divided to operate during a one hour period, the same time in each time zone). A next set of Linux hosts **10**, e.g. hosts located in the next time zone to be updated, may have policies processed according to the above example procedure.

(67) A preparation stage may perform tasks such as have an orchestrator process **70** calculate or determine those hosts expected to have policies calculated for them but did not, or hosts who were expected but did not “show up”; calculate a new host list for a current time period (e.g. the current day) by adding those hosts who were expected and had policies determined (e.g. “showed up) and in addition those hosts who did not (e.g. were unexpected) together; and publishing such a combined list of hosts onto message bus (e.g. using “topic #1: calculate”). A preparation stage, and calculating policies based on the information in the preparation stage, may be performed in advance, e.g. in the cycle time period prior to when the policies are used, e.g. 24 hours before. Queries are expected to arrive, in one embodiment, an hour before policies are distributed. These preparation stage tasks may be based on assumptions that 1) if a host needs firewall policy today, then it likely still needs that same policy in the next time period or cycle; and 2) if a host does not need firewall policies today, then it likely will not need them in the next time period or cycle. Such an embodiment may improve prior art systems, which may calculate or prepare policies for a host only in response to a query from the host, which may be time consuming. An embodiment may, instead, maintain a table (e.g. in reader service **52**) of hosts that had policies prepared in the previous cycle or time period, and use this table in the current time period to prepare policies for a host in the table in advance of that host sending a query for policies. Removing hosts which do not query in the current time period may lower compute burden for decommissioned hosts. The process may restart with a calculate process with orchestrator **70** being executed or be triggered, for example as in FIG. 2.

(68) For example, Table 3 below shows example hosts A, B, C, D and E which had rule or policy data calculated or prepared for a certain day (e.g. “today”); however only host A, B and C queried reader service **52** to refresh firewall rules; whereas hosts D and E did not do so (e.g. did not show up). In addition, in the example, hosts F and G, “appeared” and queried reader service **52** on this same day. As a result, for the next time period (e.g. tomorrow), the preparation list (e.g. stored in the memory of orchestrator **70**) may be hosts A, B, C, F and G. A corresponding SQL command for such a situation may be “SELECT hostname FROM query_record WHERE queried=1 OR is_freshman=1”.

(69) TABLE-US-00001 TABLE 3 hostname Queried is_freshman A 1 B 1 C 1 D 0 E 0 F 1 G 1

(70) A clean stage may include a process such as a microservice that routinely or periodically cleans up, e.g. removes unused data or files from, a directory or other data store that is older than a threshold time, e.g. more than seven days old. During a clean process, an embodiment may delete individual files first then remove the directory that included the deleted files; otherwise there is a chance that some process (e.g. a diff service) still holds an individual file, and as a result the directory cannot be removed immediately. For example, in the sample directory structure in Table 4 below, the general file directory architecture is date.fwdarw. (subdirectory) hostname.fwdarw.file. In such a case the deletion sequence may be 1. Delete file(s) (e.g., policy.txt etc.).fwdarw.2. Delete hostname indexed directory (e.g. /hostname_1).fwdarw.3. Delete date indexed directory (e.g. /2023-12-30)

(71) TABLE-US-00002 TABLE 4 /2023-12-30 /hostname_1 -firewall.gzip -policy.txt
-2023-12-30_2023_12-29.diff -2023-12-30_2023-12-28.diff /hostname_x -
firewall.gzip -policy.txt -2023-12-30_2023_12-29.diff -2023-12-30_2023-12-28.diff
/hostname_y -firewall.gzip -policy.txt -2023-12-30_2023_12-29.diff -2023-
12-30_2023-12-28.diff /hostname_1 -firewall.gzip -policy.txt -2023-12-
30_2023_12-29.diff -2023-12-30_2023-12-28.diff

(72) The policies relevant to a host may be indexed with a composite key, for example, a key including the hostname and other information such as the IP address for the host. In a large enterprise, hostnames may not change often but the IP address may change over time. For example, if the enterprise executes a re-organization and some projects, teams or departments move into another department, the hardware inventory, including Linux hosts, might change to a different

hardware. This may cause an IP address change of affected Linux hosts. In such an embodiment, when the Linux agent sends a request to a reader service, it may also send its IP address. The reader service may use such a composite key (e.g., hostname+IP address) to check against cache service data, to confirm whether a firewall has been calculated for this host. If yes, a result may be returned from a shared server; if not a calculation may be performed from scratch. When a calculator service calculates firewall rules for a given host, it may also use the previous IP address of that host.

(73) Reference is made to FIG. 7, showing a high-level block diagram of an exemplary computing device according to some embodiments of the present invention. Computing device **300** may include a controller **305** that may be, for example, a central processing unit processor (CPU) or any other suitable multi-purpose or specific processors or controllers, a chip or any suitable computing or computational device, an operating system **315**, a memory **320**, executable code **325**, a storage system **330**, input devices **335** and output devices **340**. Controller **305** (or one or more controllers or processors, possibly across multiple units or devices) may be configured to carry out methods described herein, and/or to execute or act as the various modules, units, etc. for example when executing code **325**. More than one computing device **300** may be included in, and one or more computing devices **300** may be, or act as the components of, a system according to embodiments of the invention. Various components, computers, and modules of FIG. 1 may be or include devices such as computing device **300**, and one or more devices such as computing device **300** may carry out functions such as those described in FIGS. 2-6.

(74) Operating system **315** may be or may include any code segment (e.g., one similar to executable code **325**) designed and/or configured to perform tasks involving coordination, scheduling, arbitration, controlling or otherwise managing operation of computing device **300**, for example, scheduling execution of software programs or enabling software programs or other modules or units to communicate.

(75) Memory **320** may be or may include, for example, a Random Access Memory (RAM), a read only memory (ROM), a Dynamic RAM (DRAM), a double data rate (DDR) memory chip, a Flash memory, a volatile memory, a non-volatile memory, a cache memory, or other suitable memory or storage units. Memory **320** may be or may include a plurality of, possibly different memory units. Memory **320** may be a computer or processor non-transitory readable medium, or a computer non-transitory storage medium, e.g., a RAM.

(76) Executable code **325** may be any executable code, e.g., an application, a program, a process, task or script. Executable code **325** may be executed by controller **305** possibly under control of operating system **315**. For example, executable code **325** may configure controller **305** to act as a calculator service, orchestrator, difference calculator, simplify services or other module, or to access, filter, distribute and/or install policies, to act as a host or a server, and perform other methods as described herein. Although, for the sake of clarity, a single item of executable code **325** is shown in FIG. 7, a system according to some embodiments of the invention may include a plurality of executable code segments similar to executable code **325** that may be loaded into memory **320** or another non-transitory storage medium and cause controller **305**, when executing code **325**, to carry out methods described herein.

(77) Storage system **330** may be or may include, for example, a hard disk drive, a CD-Recordable (CD-R) drive, a universal serial bus (USB) device or other suitable removable and/or fixed storage unit. Data such as policies may be stored in storage system **330** and may be loaded from storage system **330** into memory **320**. Some of the components shown in FIG. 7 may be omitted. Although shown as a separate component, storage system **330** may be embedded or included in memory **320**.

(78) Input devices **335** may be or may include a mouse, a keyboard, a microphone, a touch screen or pad or any suitable input device. Any suitable number of input devices may be used as shown by block **335**. Output devices **340** may include one or more displays or monitors, speakers and/or any other suitable output devices. Any suitable number of output devices may be used as shown by block **340**. Any applicable input/output (I/O) devices may be connected to computing device **300** as

shown by blocks 335 and 340. For example, a wired or wireless network interface card (NIC), a printer, a universal serial bus (USB) device or external hard drive may be included in input devices 335 and/or output devices 340.

(79) In some embodiments, device 300 may include or may be, for example, a personal computer, a desktop computer, a laptop computer, a workstation, a server computer, a network device, or any other suitable computing device.

(80) In the description and claims of the present application, each of the verbs, “comprise”, “include” and “have”, and conjugates thereof, are used to indicate that the object or objects of the verb are not necessarily a complete listing of components, elements or parts of the subject or subjects of the verb. Unless otherwise stated, adjectives such as “substantially” and “about” modifying a condition or relationship characteristic of a feature or features of an embodiment of the disclosure, are understood to mean that the condition or characteristic is defined to within tolerances that are acceptable for operation of an embodiment as described. In addition, the word “or” is considered to be the inclusive “or” rather than the exclusive or, and indicates at least one of, or any combination of items it conjoins.

(81) Descriptions of embodiments of the invention in the present application are provided by way of example and are not intended to limit the scope of the invention. The described embodiments comprise different features, not all of which are required in all embodiments. Embodiments comprising different combinations of features noted in the described embodiments, will occur to a person having ordinary skill in the art. Some elements described with respect to one embodiment may be combined with features or elements described with respect to other embodiments. The scope of the invention is limited only by the claims.

(82) While certain features of the invention have been illustrated and described herein, many modifications, substitutions, changes, and equivalents may occur to those skilled in the art. It is, therefore, to be understood that the appended claims are intended to cover all such modifications and changes as fall within the true spirit of the invention.

Claims

1. A method for updating network policies, the method comprising: for each host in a subset of hosts, determining a plurality of policies relevant to the host, wherein each policy of the plurality of policies permits or prevents a network connection between the host and one or more computer systems, and determining a difference between the plurality of policies relevant to the host and a plurality of policies relevant to the host determined at an earlier time, wherein the determining of the difference comprises, for each host in the subset of hosts, by a corresponding difference service of a plurality of difference services: generating a difference file, the difference file recording deleted policies; and storing the generated difference file in a directory storing the plurality of policies relevant to the host; wherein the difference service for the host consumes files to compare from a message bus, the files to compare output by a corresponding simplify service for the host, the simplify service for the host filtering information from a compressed policy file, the compressed policy file comprising the plurality of policies relevant to the host; and for each host in the subset of hosts, installing the plurality of policies relevant to the host.
2. The method of claim 1, wherein the difference is determined by comparing a file containing the plurality of policies relevant to the host, and a file containing the plurality of policies relevant to the host determined at an earlier time.
3. The method of claim 1, wherein: the difference is saved in a first file; and the plurality of policies relevant to the host is saved in a second file; and the first file and the second file are stored in the same directory.
4. The method of claim 3, wherein the directory is indexed by the host relevant to the policies and the time period relevant to the policies.

5. The method of claim 1 comprising: determining, from a set of hosts, a subset of hosts to have policies updated, the determining based on a table; adding a host to the table if a query is received from the host; and removing a host from the table if a query is not received from the host during a predetermined time period.
6. The method of claim 1, wherein the subset of hosts is determined based on the time zone corresponding to the location of each host in the subset of hosts.
7. The method of claim 1, wherein a policy relevant to a host is a policy that is designed to work with the host.
8. A system for updating network policies, the system comprising: at least one memory; and a plurality of processors, wherein: at least one processor among the plurality of processors is configured to for each host in the subset of hosts, determine a plurality of policies relevant to the host, wherein each policy of the plurality of policies permits or prevents a network connection between the host and one or more computer systems, and determine a difference between the plurality of policies relevant to the host and a plurality of policies relevant to the host determined at an earlier time, wherein the determining of the difference comprises, for each host in the subset of hosts, by a corresponding difference service of a plurality of difference services: generating a difference file, the difference file recording deleted policies; and storing the generated difference file in a directory storing the plurality of policies relevant to the host; wherein the difference service for the host consumes files to compare from a message bus, the files to compare output by a corresponding simplify service for the host, the simplify service for the host filtering information from a compressed policy file, the compressed policy file comprising the plurality of policies relevant to the host; and for each host in the subset of hosts, at least one processor among the plurality of processors is configured to install the plurality of policies relevant to the host.
9. The system of claim 8, wherein the difference is determined by comparing a file containing the plurality of policies relevant to the host, and a file containing the plurality of policies relevant to the host determined at an earlier time.
10. The system of claim 8, wherein: the difference is saved in a first file; and the plurality of policies relevant to the host is saved in a second file; and the first file and the second file are stored in the same directory.
11. The system of claim 10, wherein the directory is indexed by the host relevant to the policies and the time period relevant to the policies.
12. The system of claim 8 wherein at least one processor among the plurality of processors is configured to: determine, from a set of hosts, a subset of hosts to have policies updated, the determining based on a table; add a host to the table if a query is received from the host; and remove a host from the table if a query is not received from the host during a predetermined time period.
13. The system of claim 8, wherein the subset of hosts is determined based on the time zone corresponding to the location of each host in the subset of hosts.
14. The system of claim 8, wherein a policy relevant to a host is a policy that is designed to work with the host.
15. A method for distributing network rules, the method comprising: determining for each of a set of hosts, a plurality of rules to be applied to the host, wherein each rule of the plurality of rules permits or prevents a network connection between the host and one or more computer systems, and determining a difference between the plurality of rules to be applied to the host and a previous plurality of rules to be applied to the host, wherein the determining of the difference comprises, for each host in the set of hosts, by a corresponding difference service of a plurality of difference services: generating a difference file, the difference file recording deleted rules; and storing the generated difference file in a directory storing the plurality of rules to be applied to the host; wherein the difference service for the host consumes files to compare from a message bus, the files to compare output by a corresponding simplify service for the host, the simplify service for the host

filtering information from a compressed rule file, the compressed rule file comprising the plurality of rules to be applied to the host; and for each host of the set of hosts, installing the plurality of rules to be applied to the host.

16. The method of claim 15, wherein the difference is determined by comparing a file containing the plurality of rules to be applied to the host, and a previous file of rules to be applied to the host.

17. The method of claim 15, wherein: the difference is saved in a first file; and the plurality of rules to be applied to the host is saved in a second file; and the first file and the second file are stored in the same directory.

18. The method of claim 17, wherein the directory is indexed by the host relevant to the rules and the time period relevant to the rules.

19. The method of claim 15 comprising: determining, from a set of hosts, a subset of hosts to have rules updated, the determining based on a table; adding a host to the table if a query is received from the host; and removing a host from the table if a query is not received from the host during a predetermined time period.

20. The method of claim 15, wherein a rule to be applied to a host is a rule that is designed to work with the host.
