



US012407701B1

(12) **United States Patent**
Varakantam et al.

(10) **Patent No.:** **US 12,407,701 B1**
(45) **Date of Patent:** **Sep. 2, 2025**

(54) **COMMUNITY-BASED GENERATION OF
POLICIES FOR A DATA PLATFORM**

(71) Applicant: **Lacework, Inc.**, San Jose, CA (US)

(72) Inventors: **Vikram R. Varakantam**, Fremont, CA (US); **Yijou Chen**, Cupertino, CA (US); **Bharath Kolar Venkat**, San Jose, CA (US); **George B. Spofford**, Redwood City, CA (US); **Mandar Pimpale**, San Jose, CA (US)

(73) Assignee: **Fortinet, Inc.**, Sunnyvale, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 457 days.

(21) Appl. No.: **17/589,770**

(22) Filed: **Jan. 31, 2022**

Related U.S. Application Data

(63) Continuation-in-part of application No. 17/510,179, filed on Oct. 25, 2021, now Pat. No. 11,637,849, (Continued)

(51) **Int. Cl.**
H04L 9/40 (2022.01)
G06F 9/455 (2018.01)
(Continued)

(52) **U.S. Cl.**
CPC **H04L 63/1425** (2013.01); **G06F 9/455** (2013.01); **G06F 9/545** (2013.01);
(Continued)

(58) **Field of Classification Search**
CPC . H04L 63/1425; H04L 67/535; H04L 43/045; H04L 43/06; H04L 63/10;
(Continued)

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,584,024 A 12/1996 Shwartz
5,806,062 A 9/1998 Chen
(Continued)

FOREIGN PATENT DOCUMENTS

EP 3489828 A1 5/2019
KR 101039698 12/2010
(Continued)

OTHER PUBLICATIONS

Dierks et al. "On the cluster admission problem for cloud computing", arXiv:1804.07571, Apr. 20, 2018. (Year: 2018).*
(Continued)

Primary Examiner — Peter C Shaw


(74) *Attorney, Agent, or Firm* — Jaffery Watson Hamilton & DeSanctis LLP

(57) **ABSTRACT**

A disclosed data platform may be configured to access a custom policy developed by a particular client entity that uses the data platform. The custom policy may invoke a query that targets a target dataset ingested from a cloud environment and stored in a data store. The data platform may determine that the custom policy is likely to be of value to one or more other client entities, besides the particular client entity, that also use the data platform. In response to the determining that the custom policy is likely to be of value to the one or more other client entities, the data platform may generate a community policy based on the custom policy. The community policy may be available for use by the one or more other client entities. Corresponding methods, systems, and products are also disclosed.

19 Claims, 60 Drawing Sheets

700



Related U.S. Application Data

- which is a continuation of application No. 16/786,822, filed on Feb. 10, 2020, now Pat. No. 11,157,502, which is a continuation of application No. 16/134,806, filed on Sep. 18, 2018, now Pat. No. 10,614,071.
- (60) Provisional application No. 63/248,915, filed on Sep. 27, 2021, provisional application No. 63/239,275, filed on Aug. 31, 2021, provisional application No. 63/228,980, filed on Aug. 3, 2021, provisional application No. 63/224,373, filed on Jul. 21, 2021, provisional application No. 62/650,971, filed on Mar. 30, 2018, provisional application No. 62/590,986, filed on Nov. 27, 2017.

(51) Int. Cl.

G06F 9/54 (2006.01)
G06F 16/2455 (2019.01)
G06F 16/901 (2019.01)
G06F 16/9038 (2019.01)
G06F 16/9535 (2019.01)
G06F 16/9537 (2019.01)
G06F 21/57 (2013.01)
H04L 43/045 (2022.01)
H04L 43/06 (2022.01)
H04L 67/306 (2022.01)
H04L 67/50 (2022.01)

(52) U.S. Cl.

CPC **G06F 16/9024** (2019.01); **G06F 16/9038** (2019.01); **G06F 16/9535** (2019.01); **G06F 16/9537** (2019.01); **G06F 21/57** (2013.01); **H04L 43/045** (2013.01); **H04L 43/06** (2013.01); **H04L 63/10** (2013.01); **H04L 67/306** (2013.01); **H04L 67/535** (2022.05); **G06F 16/2456** (2019.01)

(58) Field of Classification Search

CPC H04L 67/306; G06F 16/9038; G06F 16/9535; G06F 16/9537; G06F 16/9024; G06F 9/455; G06F 9/545; G06F 21/57; G06F 16/2456
 USPC 707/770
 See application file for complete search history.

(56) References Cited**U.S. PATENT DOCUMENTS**

6,347,339 B1 2/2002 Morris et al.
 6,363,411 B1 3/2002 Dugan et al.
 6,434,663 B1 8/2002 Grimsrud et al.
 6,507,852 B1 1/2003 Dempsey et al.
 6,938,084 B2 8/2005 Gamache et al.
 7,054,873 B2 5/2006 Nordström et al.
 7,233,333 B2 6/2007 Lomask
 7,310,733 B1 12/2007 Pearson et al.
 7,478,246 B2 1/2009 Arndt et al.
 7,484,091 B2 1/2009 Bade et al.
 7,526,501 B2 4/2009 Albahari et al.
 7,529,801 B2 5/2009 Moore et al.
 7,558,883 B1* 7/2009 Lampion G06F 16/27
 7,562,045 B2 7/2009 Beadle et al.
 7,707,411 B2 4/2010 Bade et al.
 7,739,211 B2 6/2010 Coffman et al.
 7,743,153 B2 6/2010 Hall et al.
 7,747,559 B2 6/2010 Leitner et al.
 7,765,431 B2 7/2010 Agha et al.
 7,797,548 B2 9/2010 Pearson et al.
 7,856,544 B2 12/2010 Schenfeld et al.
 7,926,026 B2 4/2011 Klein et al.

7,962,635 B2 6/2011 Naidu et al.
 7,996,885 B2 8/2011 Jaiswal et al.
 8,032,925 B2 10/2011 Cho
 8,037,284 B2 10/2011 Schenfeld et al.
 8,037,521 B2 10/2011 Minato
 8,050,907 B2 11/2011 Baisley et al.
 8,086,852 B2 12/2011 Bade et al.
 8,103,906 B1 1/2012 Alibakhsh
 8,122,122 B1 2/2012 Clingenpeel
 8,140,977 B2 3/2012 Kriss et al.
 8,151,107 B2 4/2012 Song et al.
 8,160,999 B2 4/2012 Jin et al.
 8,209,204 B2 6/2012 Adler et al.
 8,276,197 B1 9/2012 Mangal et al.
 8,291,233 B2 10/2012 Pearson et al.
 8,301,660 B2 10/2012 Yalamanchi
 8,341,711 B1 12/2012 Pennington et al.
 8,351,456 B2 1/2013 Kadous et al.
 8,352,589 B2 1/2013 Ridel et al.
 8,359,584 B2 1/2013 Rao et al.
 8,443,442 B2 5/2013 Wang
 8,490,055 B2 7/2013 Basak
 8,497,863 B2 7/2013 Xie et al.
 8,549,002 B2 10/2013 Herter et al.
 8,561,157 B2 10/2013 Ge
 8,595,262 B1 11/2013 Hayden
 8,607,306 B1 12/2013 Bridge et al.
 8,655,989 B2 2/2014 Ritter et al.
 8,671,453 B2 3/2014 Underwood
 8,725,587 B2 5/2014 Beadle et al.
 8,826,403 B2 9/2014 Bhaskaran et al.
 8,843,646 B2 9/2014 Kuzin et al.
 8,862,524 B2 10/2014 Zheng
 8,959,608 B2 2/2015 Ahmed et al.
 9,021,583 B2 4/2015 Wittenstein
 9,037,273 B2 5/2015 Mikkelsen
 9,043,764 B2 5/2015 Ranganathan et al.
 9,053,306 B2 6/2015 Yoshigaki et al.
 9,053,437 B2 6/2015 Adler et al.
 9,064,210 B1 6/2015 Hart
 9,075,618 B2 7/2015 Winternitz et al.
 9,110,873 B2 8/2015 Woodall et al.
 9,159,024 B2 10/2015 Bhanot et al.
 9,189,623 B1 11/2015 Lin et al.
 9,225,730 B1 12/2015 Brezinski
 9,231,935 B1 1/2016 Bridge et al.
 9,239,873 B2 1/2016 Branch et al.
 9,246,897 B2 1/2016 He
 9,323,806 B2 4/2016 Sadikov et al.
 9,332,020 B2 5/2016 Thomas
 9,369,450 B1 6/2016 Barak et al.
 9,391,978 B2 7/2016 Burch et al.
 9,400,882 B2 7/2016 Pearson et al.
 9,430,830 B2 8/2016 Madabhushi et al.
 9,495,522 B2 11/2016 Singh et al.
 9,497,224 B2 11/2016 Sweet et al.
 9,515,999 B2 12/2016 Ylonen
 9,516,053 B1 12/2016 Muddu et al.
 9,537,851 B2 1/2017 Gordon et al.
 9,558,265 B1 1/2017 Tacchi
 9,569,869 B2 2/2017 Hesse et al.
 9,582,766 B2 2/2017 Sadikov et al.
 9,589,069 B2 3/2017 Yang et al.
 9,591,010 B1 3/2017 Muddu et al.
 9,596,253 B2 3/2017 Chauhan
 9,596,254 B1 3/2017 Muddu et al.
 9,596,295 B2 3/2017 Banadaki et al.
 9,600,915 B2 3/2017 Winternitz et al.
 9,602,506 B2 3/2017 Kang et al.
 9,602,526 B2 3/2017 Liu et al.
 9,639,676 B2 5/2017 Betz et al.
 9,652,875 B2 5/2017 Vassilvitskii et al.
 9,654,503 B1 5/2017 Kowalyshyn
 9,659,337 B2 5/2017 Lee et al.
 9,665,660 B2 5/2017 Wensel
 9,667,641 B2 5/2017 Muddu et al.
 9,679,243 B2 6/2017 Zou et al.
 9,699,205 B2 7/2017 Muddu et al.
 9,710,332 B1 7/2017 Fan et al.

(56)

References Cited

U.S. PATENT DOCUMENTS

9,720,703 B2	8/2017	Reick et al.	10,496,263 B2	12/2019	So et al.
9,720,704 B2	8/2017	Reick et al.	10,496,468 B2	12/2019	Gefen et al.
9,727,441 B2	8/2017	Agarwal et al.	10,496,678 B1	12/2019	Tang
9,727,604 B2	8/2017	Jin et al.	10,505,818 B1	12/2019	Yona et al.
9,729,416 B1	8/2017	Khanal et al.	10,505,990 B1 *	12/2019	Natarajan H04L 63/205
9,734,040 B2	8/2017	Gounares	10,510,007 B2	12/2019	Singhal et al.
9,740,744 B2	8/2017	Stetson et al.	10,515,095 B2	12/2019	Childress et al.
9,741,138 B2	8/2017	Friedlander et al.	10,516,669 B2 *	12/2019	Bourgeois G06F 16/178
9,749,339 B2	8/2017	Kadambe	10,521,584 B1	12/2019	Mehr
9,753,960 B1	9/2017	Troyanovsky	10,534,633 B2	1/2020	Hilemon et al.
9,760,619 B1	9/2017	Lattanzi et al.	10,565,373 B1	2/2020	Rao et al.
9,781,115 B2	10/2017	Heise	10,581,891 B1	3/2020	Kapoor et al.
9,787,705 B1	10/2017	Love et al.	10,587,609 B2	3/2020	Ebrahimi et al.
9,805,080 B2	10/2017	Joshi et al.	10,592,535 B2	3/2020	Ahn et al.
9,805,140 B2	10/2017	Chakrabarti et al.	10,594,718 B1	3/2020	Deaguero
9,811,790 B2	11/2017	Ahern et al.	10,599,718 B2	3/2020	Kumar et al.
9,813,285 B1	11/2017	McGovern et al.	RE47,937 E	4/2020	Ramachandran et al.
9,813,435 B2	11/2017	Muddu et al.	RE47,952 E	4/2020	Ramachandran et al.
9,819,671 B2	11/2017	Ji	10,614,200 B2	4/2020	Betz et al.
9,824,473 B2	11/2017	Winternitz et al.	10,642,867 B2	5/2020	Palanciuc
9,830,435 B2	11/2017	Haven	10,656,979 B2	5/2020	Ishakian et al.
9,836,183 B1	12/2017	Love et al.	10,664,757 B2	5/2020	Lastras-Montano et al.
9,838,410 B2	12/2017	Muddu et al.	10,666,668 B2	5/2020	Muddu et al.
9,843,837 B2	12/2017	Gopalan	10,673,880 B1	6/2020	Pratt et al.
9,852,230 B2	12/2017	Fleury et al.	10,685,295 B1	6/2020	Ross et al.
9,853,968 B2	12/2017	Shen	10,693,900 B2	6/2020	Zadeh et al.
9,864,672 B2	1/2018	Seto et al.	10,698,954 B2	6/2020	Piechowicz et al.
9,887,999 B2	2/2018	Dong et al.	10,701,051 B2	6/2020	Ohsumi
9,923,911 B2	3/2018	Vasseur et al.	10,708,082 B1	7/2020	Bakiaraj
9,942,220 B2	4/2018	Bajenov et al.	10,735,329 B2	8/2020	Wang
9,946,800 B2	4/2018	Qian et al.	10,754,940 B2	8/2020	Ohsumi
9,953,014 B1	4/2018	Reshadi	10,756,982 B2	8/2020	Bai et al.
9,954,842 B2	4/2018	Huang	10,771,488 B2	9/2020	Verma et al.
9,985,827 B2	5/2018	Li	10,775,183 B2	9/2020	Ho et al.
10,003,605 B2	6/2018	Muddu et al.	10,776,191 B2	9/2020	Zheng et al.
10,033,611 B1	7/2018	Linkous	10,788,570 B2	9/2020	Wilson
10,104,071 B2	10/2018	Gordon et al.	10,791,131 B2	9/2020	Nor et al.
10,115,111 B2	10/2018	Miltonberger	10,797,974 B2	10/2020	Giura et al.
10,116,670 B2	10/2018	Muddu et al.	10,812,497 B2	10/2020	Venkatramani et al.
10,121,000 B1	11/2018	Rivlin et al.	10,824,675 B2	11/2020	Alonso et al.
10,122,740 B1	11/2018	Finkelshtein et al.	10,824,813 B2	11/2020	Smith et al.
10,127,273 B2	11/2018	Dickey	10,885,452 B1	1/2021	Garg
10,142,357 B1	11/2018	Tamersoy	10,904,007 B2	1/2021	Kim et al.
10,148,677 B2	12/2018	Muddu et al.	10,904,270 B2	1/2021	Muddu et al.
10,149,148 B2	12/2018	Zha et al.	10,911,470 B2	2/2021	Muddu et al.
10,158,652 B2	12/2018	Muddu et al.	10,986,114 B1	4/2021	Singh et al.
10,182,058 B2	1/2019	Xu	11,036,716 B2	6/2021	Griffith et al.
10,205,735 B2	2/2019	Apostolopoulos	11,036,800 B1	6/2021	Kayyoor et al.
10,205,736 B2	2/2019	Rieke	11,044,264 B2	6/2021	Durairaj et al.
10,237,254 B2	3/2019	McDowell et al.	11,048,492 B2	6/2021	Jain et al.
10,237,294 B1	3/2019	Zadeh et al.	11,082,289 B2	8/2021	Dang et al.
10,243,970 B2	3/2019	Muddu et al.	11,120,343 B2	9/2021	Das et al.
10,249,266 B2	4/2019	Zamir	11,126,533 B2	9/2021	Knowles et al.
10,254,848 B2	4/2019	Winternitz et al.	11,194,849 B2	12/2021	Lassoued et al.
10,331,659 B2	6/2019	Ahuja et al.	11,212,299 B2	12/2021	Gamble et al.
10,338,895 B2	7/2019	Zhang et al.	11,258,807 B2	2/2022	Muddu et al.
10,339,309 B1	7/2019	Kling	11,281,519 B2	3/2022	Krishnaswamy et al.
10,367,704 B2	7/2019	Giura et al.	11,314,789 B2	4/2022	Goldfarb
10,382,303 B2	8/2019	Khanal et al.	11,411,966 B2	8/2022	Muddu et al.
10,382,529 B2	8/2019	Wan et al.	11,431,735 B2	8/2022	Shua
10,389,738 B2	8/2019	Muddu et al.	11,463,464 B2	10/2022	Zadeh et al.
10,389,742 B2	8/2019	Devi	11,489,863 B1	11/2022	Shua
10,419,463 B2	9/2019	Muddu et al.	11,494,787 B2	11/2022	Erickson et al.
10,419,465 B2	9/2019	Muddu et al.	11,544,138 B2	1/2023	Kapish et al.
10,419,468 B2	9/2019	Glatfelter et al.	11,575,693 B1	2/2023	Muddu et al.
10,419,469 B1	9/2019	Singh et al.	11,606,272 B1	3/2023	Popelka et al.
10,425,437 B1	9/2019	Bog et al.	11,636,090 B2	4/2023	Li et al.
10,432,639 B1	10/2019	Bebec et al.	11,640,388 B2	5/2023	Yang et al.
10,432,669 B1 *	10/2019	Badhwar H04L 63/1408	11,647,034 B2	5/2023	Levin et al.
10,447,526 B2	10/2019	Tucker et al.	11,658,990 B2	5/2023	Shapoury
10,454,753 B2	10/2019	Sasturkar et al.	11,669,571 B2	6/2023	Binkley et al.
10,454,889 B2	10/2019	Huang	11,693,958 B1	7/2023	Steiman
10,459,979 B2	10/2019	Piechowicz et al.	11,722,554 B2	8/2023	Keren et al.
10,462,169 B2	10/2019	Durairaj et al.	11,734,351 B2	8/2023	Binkley et al.
10,491,705 B2	11/2019	Oetting et al.	11,734,419 B1	8/2023	Mackle
			11,748,473 B2	9/2023	Araujo et al.
			11,755,576 B1	9/2023	Jiang et al.
			11,755,602 B2	9/2023	Smith et al.
			11,769,098 B2	9/2023	Adinarayan et al.

US 12,407,701 B1

Page 4

(56)	References Cited			2011/0066841	A1 *	3/2011	Goodrow	H04L 43/0894
	U.S. PATENT DOCUMENTS			2011/0119100	A1	5/2011	Ruhl et al.	714/4.2
11,770,387	B1	9/2023	Shivamoggi et al.	2011/0154287	A1	6/2011	Mukkamala et al.	
2002/0059531	A1	5/2002	Lai	2011/0302631	A1	12/2011	Sureshchandra et al.	
2002/0161889	A1	10/2002	Gamache et al.	2011/0314148	A1	12/2011	Petersen et al.	
2002/0184225	A1	12/2002	Ghukasyan	2012/0005243	A1	1/2012	Van De Merwe	
2003/0037136	A1	2/2003	Labovitz	2012/0054732	A1	3/2012	Jain et al.	
2003/0179227	A1	9/2003	Ahmad	2012/0089875	A1	4/2012	Faust et al.	
2003/0233361	A1	12/2003	Cady	2012/0102029	A1	4/2012	Larson et al.	
2004/0015470	A1	1/2004	Smith	2012/0102543	A1 *	4/2012	Kohli	H04L 63/20
2004/0225929	A1	11/2004	Agha et al.					726/1
2005/0060287	A1	3/2005	Hellman et al.	2012/0143898	A1	6/2012	Bruno et al.	
2005/0102284	A1	5/2005	Srinivasan	2012/0158858	A1	6/2012	Gkantsidis et al.	
2005/0102365	A1	5/2005	Moore et al.	2012/0159333	A1	6/2012	Mital et al.	
2005/0108142	A1	5/2005	Beadle et al.	2012/0173541	A1	7/2012	Venkataramani	
2005/0188222	A1	8/2005	Motsinger	2012/0221535	A1	8/2012	Dubbels et al.	
2005/0222996	A1	10/2005	Yalamanchi	2012/0317149	A1	12/2012	Jagota et al.	
2005/0231760	A1	10/2005	Minato	2012/0317151	A1	12/2012	Ruf	
2005/0246288	A1	11/2005	Kimura et al.	2012/0323956	A1	12/2012	Dumitru	
2005/0246521	A1	11/2005	Bade et al.	2013/0024412	A1	1/2013	Gong et al.	
2006/0025987	A1	2/2006	Baisley et al.	2013/0067100	A1	3/2013	Kuzin et al.	
2006/0026419	A1	2/2006	Arndt et al.	2013/0081118	A1	3/2013	Ge	
2006/0036896	A1	2/2006	Gamache et al.	2013/0086667	A1	4/2013	Haven	
2006/0085437	A1	4/2006	Brodhun	2013/0097320	A1	4/2013	Ritter et al.	
2006/0090095	A1	4/2006	Massa et al.	2013/0151453	A1	6/2013	Bhanot et al.	
2006/0109271	A1	5/2006	Lomask	2013/0173915	A1	7/2013	Haulund	
2006/0259470	A1	11/2006	Chandrasekharan	2013/0219295	A1	8/2013	Feldman et al.	
2006/0288415	A1	12/2006	Wong	2013/0269007	A1	10/2013	Yoshigaki et al.	
2007/0050497	A1	3/2007	Haley	2013/0304915	A1	11/2013	Kawai	
2007/0118909	A1	5/2007	Hertzog	2013/0305357	A1	11/2013	Ayyagari	
2007/0130330	A1	6/2007	Ridel et al.	2013/0325890	A1	12/2013	Vaught	
2007/0162605	A1	7/2007	Chalasani et al.	2014/0041005	A1	2/2014	He	
2007/0162963	A1	7/2007	Penet et al.	2014/0067750	A1	3/2014	Ranganathan et al.	
2007/0168696	A1	7/2007	Ridel et al.	2014/0098101	A1	4/2014	Friedlander et al.	
2007/0169175	A1	7/2007	Hall et al.	2014/0115001	A1	4/2014	Arroyo	
2007/0214111	A1	9/2007	Jin et al.	2014/0115011	A1	4/2014	Buerner	
2007/0225956	A1	9/2007	Pratt et al.	2014/0125672	A1	5/2014	Winternitz et al.	
2007/0266425	A1	11/2007	Cho	2014/0149148	A1	5/2014	Luciani	
2007/0282916	A1	12/2007	Albahari et al.	2014/0181944	A1	6/2014	Ahmed et al.	
2008/0034411	A1	2/2008	Aoyama	2014/0208191	A1	7/2014	Zaric	
2008/0065879	A1	3/2008	Song et al.	2014/0211988	A1	7/2014	Fan et al.	
2008/0072062	A1	3/2008	Pearson et al.	2014/0245443	A1	8/2014	Chakraborty	
2008/0109730	A1	5/2008	Coffman et al.	2014/0279779	A1	9/2014	Zou et al.	
2008/0147707	A1	6/2008	Jin et al.	2014/0280068	A1	9/2014	Dhoopar	
2008/0148180	A1	6/2008	Liu	2014/0325631	A1	10/2014	Pearson et al.	
2008/0151893	A1	6/2008	Nordmark	2014/0359558	A1	12/2014	Chamberlain	
2008/0155335	A1	6/2008	Klein et al.	2014/0379716	A1	12/2014	Branch et al.	
2008/0184201	A1 *	7/2008	Burns	2015/0058619	A1	2/2015	Sweet et al.	
			H04L 41/0894	2015/0135312	A1	5/2015	Wada	
			717/121	2015/0161201	A1	6/2015	Sadikov et al.	
2008/0244718	A1	10/2008	Frost et al.	2015/0188751	A1	7/2015	Vasseur	
2008/0263643	A1	10/2008	Jaiswal et al.	2015/0193523	A1 *	7/2015	Cox	G06N 5/027
2008/0270451	A1	10/2008	Thomsen et al.					707/737
2009/0006843	A1	1/2009	Bade et al.	2015/0213598	A1	7/2015	Madabhushi et al.	
2009/0007010	A1	1/2009	Kriss et al.	2015/0310649	A1	10/2015	Winternitz et al.	
2009/0019160	A1	1/2009	Schuler	2015/0341379	A1	11/2015	Lefebvre	
2009/0063857	A1	3/2009	Bade et al.	2015/0356144	A1	12/2015	Chawla	
2009/0165109	A1	6/2009	Hird	2016/0063226	A1	3/2016	Singh et al.	
2009/0177573	A1	7/2009	Beadle et al.	2016/0078365	A1	3/2016	Baumard	
2009/0222740	A1	9/2009	Yuan	2016/0080404	A1	3/2016	Kohout	
2009/0228474	A1	9/2009	Chiu	2016/0110434	A1	4/2016	Kakaraddi	
2009/0271504	A1	10/2009	Ginter	2016/0120070	A1	4/2016	Myrah et al.	
2009/0287720	A1	11/2009	Herter et al.	2016/0149937	A1	5/2016	Katmor	
2009/0307651	A1	12/2009	Senthil et al.	2016/0203411	A1	7/2016	Sadikov et al.	
2009/0327328	A1	12/2009	Woodall et al.	2016/0205125	A1	7/2016	Kim	
2010/0042823	A1	2/2010	Arndt et al.	2016/0218911	A1	7/2016	Wessels	
2010/0094767	A1	4/2010	Miltonberger	2016/0261544	A1	9/2016	Conover	
2010/0114931	A1	5/2010	Xie	2016/0275303	A1	9/2016	Narayanawamy et al.	
2010/0122270	A1	5/2010	Lin	2016/0330183	A1	11/2016	McDowell et al.	
2010/0172261	A1	7/2010	Shinbo	2016/0330206	A1	11/2016	Xu	
2010/0217860	A1	8/2010	Naidu et al.	2016/0357521	A1	12/2016	Zhang et al.	
2010/0274785	A1	10/2010	Procopiu	2016/0359592	A1	12/2016	Kulshreshtha	
2010/0309206	A1	12/2010	Xie et al.	2016/0359872	A1	12/2016	Yadav	
2010/0312740	A1	12/2010	Clemm et al.	2016/0373428	A1	12/2016	Shi	
2010/0329162	A1	12/2010	Kadous et al.	2017/0063830	A1	3/2017	Huang	
2011/0023098	A1	1/2011	Pearson et al.	2017/0063888	A1	3/2017	Muddu et al.	
2011/0029952	A1	2/2011	Harrington	2017/0063903	A1	3/2017	Muddu et al.	
2011/0055138	A1	3/2011	Khanduja	2017/0063905	A1	3/2017	Muddu et al.	

(56)	References Cited		2018/0262592 A1*	9/2018	Zandi	H04L 67/55
	U.S. PATENT DOCUMENTS		2018/0268078 A1	9/2018	Gianetto	
			2018/0288063 A1	10/2018	Koottayi	
			2018/0314740 A1	11/2018	Van Osten et al.	
2017/0063906 A1	3/2017	Muddu et al.	2018/0329958 A1	11/2018	Choudhury	
2017/0063908 A1	3/2017	Muddu et al.	2018/0336353 A1	11/2018	Manadhata	
2017/0063909 A1	3/2017	Muddu et al.	2018/0357422 A1	12/2018	Telang	
2017/0063910 A1	3/2017	Muddu et al.	2018/0359162 A1	12/2018	Savov et al.	
2017/0063911 A1	3/2017	Muddu et al.	2018/0367548 A1	12/2018	Stokes, III	
2017/0063912 A1	3/2017	Muddu et al.	2019/0028327 A1	1/2019	Silva	
2017/0070594 A1	3/2017	Oetting et al.	2019/0042727 A1	2/2019	Rodniansky	
2017/0076206 A1	3/2017	Lastras-Montano et al.	2019/0042879 A1	2/2019	Munoz	
2017/0085553 A1	3/2017	Gordon et al.	2019/0042950 A1	2/2019	Lin et al.	
2017/0086069 A1	3/2017	Liu	2019/0050445 A1	2/2019	Griffith et al.	
2017/0102961 A1	4/2017	Hilemon et al.	2019/0058626 A1	2/2019	Knowles et al.	
2017/0111245 A1	4/2017	Ishakian et al.	2019/0068598 A1	2/2019	Kang et al.	
2017/0116315 A1	4/2017	Xiong et al.	2019/0075126 A1	3/2019	Muddu et al.	
2017/0118099 A1	4/2017	Huang	2019/0087480 A1	3/2019	Palanciuc	
2017/0118240 A1	4/2017	Devi Reddy	2019/0101622 A1	4/2019	Wilson	
2017/0134240 A1	5/2017	Hevizi	2019/0109870 A1	4/2019	Bedhapudi	
2017/0142140 A1	5/2017	Muddu et al.	2019/0132224 A1	5/2019	Verma	
2017/0147646 A1	5/2017	Lee	2019/0147180 A1	5/2019	Rodniansky et al.	
2017/0148197 A1	5/2017	Winternitz et al.	2019/0149553 A1	5/2019	Xu	
2017/0155570 A1	6/2017	Maheshwari et al.	2019/0158524 A1	5/2019	Zadeh et al.	
2017/0155672 A1	6/2017	Muthukrishnan et al.	2019/0163555 A1	5/2019	Zheng et al.	
2017/0163666 A1	6/2017	Venkatramani	2019/0222597 A1	7/2019	Crabtree	
2017/0208050 A1	7/2017	Iasi et al.	2019/0227860 A1	7/2019	Gefen et al.	
2017/0223036 A1	8/2017	Muddu et al.	2019/0259033 A1	8/2019	Reddy	
2017/0230183 A1	8/2017	Sweet et al.	2019/0289058 A1*	9/2019	Bhoj	H04L 67/535
2017/0249069 A1	8/2017	Zamir	2019/0312796 A1	10/2019	Giura et al.	
2017/0251013 A1	8/2017	Kirti	2019/0312898 A1	10/2019	Verma et al.	
2017/0257358 A1	9/2017	Ebrahimi et al.	2019/0327251 A1	10/2019	Muddu et al.	
2017/0262521 A1	9/2017	Cho et al.	2019/0339965 A1	11/2019	Garvey	
2017/0272344 A1	9/2017	Tang	2019/0342282 A1	11/2019	Carbune et al.	
2017/0277553 A1	9/2017	Zada et al.	2019/0342307 A1	11/2019	Gamble et al.	
2017/0277997 A1	9/2017	Zong et al.	2019/0342311 A1	11/2019	Muddu et al.	
2017/0279827 A1	9/2017	Savalle	2019/0349305 A1	11/2019	Wang	
2017/0286190 A1	10/2017	Ishakian et al.	2019/0354554 A1	11/2019	Piechowicz et al.	
2017/0288974 A1	10/2017	Yoshihira	2019/0356555 A1	11/2019	Bai et al.	
2017/0318027 A1	11/2017	Rodniansky et al.	2019/0364067 A1	11/2019	Yona et al.	
2017/0330096 A1	11/2017	Gupta et al.	2019/0378050 A1	12/2019	Edkin	
2017/0337262 A1	11/2017	Smith et al.	2019/0391964 A1*	12/2019	Studer	G06Q 30/02
2017/0346683 A1	11/2017	Li et al.	2020/0014718 A1	1/2020	Durairaj et al.	
2017/0353853 A1	12/2017	Zha et al.	2020/0021607 A1	1/2020	Muddu et al.	
2017/0359361 A1	12/2017	Modani	2020/0065857 A1	2/2020	Lagi et al.	
2017/0366492 A1	12/2017	Ho	2020/0074341 A1	3/2020	He et al.	
2018/0004835 A1	1/2018	Piechowicz et al.	2020/0076685 A1	3/2020	Vaidya	
2018/0004859 A1	1/2018	Piechowicz et al.	2020/0080856 A1	3/2020	Ho et al.	
2018/0007145 A1	1/2018	Piechowicz et al.	2020/0175042 A1	6/2020	Batruni	
2018/0013650 A1	1/2018	Khanal et al.	2020/0175361 A1	6/2020	Che et al.	
2018/0019932 A1	1/2018	Giura et al.	2020/0192690 A1	6/2020	Gupta	
2018/0020015 A1	1/2018	Munro	2020/0228555 A1	7/2020	Wittenschlaeger	
2018/0025361 A1	1/2018	Ferrarons Llagostera et al.	2020/0252376 A1	8/2020	Feng	
2018/0034840 A1	2/2018	Marquardt	2020/0272740 A1	8/2020	Obee	
2018/0039688 A1	2/2018	Ahn et al.	2020/0278892 A1	9/2020	Nainar	
2018/0063178 A1	3/2018	Jadhav	2020/0280592 A1	9/2020	Ithal	
2018/0067981 A1	3/2018	Ahuja et al.	2020/0287923 A1	9/2020	Raghavendra et al.	
2018/0069885 A1	3/2018	Patterson et al.	2020/0287927 A1	9/2020	Zadeh et al.	
2018/0084069 A1	3/2018	Be'Ery et al.	2020/0311644 A1	10/2020	Willard	
2018/0089132 A1	3/2018	Atta et al.	2020/0320106 A1	10/2020	Goldfarb	
2018/0089269 A1*	3/2018	Pal	2020/0334293 A1	10/2020	Piechowicz et al.	
2018/0096047 A1	4/2018	Childress et al.	2020/0351151 A1	11/2020	Dang et al.	
2018/0097793 A1	4/2018	Agarwal et al.	2020/0403860 A1	12/2020	Lewis	
2018/0101779 A1	4/2018	Canim et al.	2020/0404008 A1	12/2020	Venkatramani et al.	
2018/0103052 A1	4/2018	Choudhury	2020/0412752 A1	12/2020	Shapoury	
2018/0115578 A1	4/2018	Subbarayan	2021/0019209 A1	1/2021	Krishnaswamy et al.	
2018/0123864 A1	5/2018	Tucker et al.	2021/0232420 A1	7/2021	Dhruvakumar et al.	
2018/0137858 A1	5/2018	Saxena	2021/0286798 A1	9/2021	Li et al.	
2018/0139200 A1	5/2018	Gordon et al.	2021/0294798 A1	9/2021	Binkley et al.	
2018/0173789 A1	6/2018	Llagostera	2021/0329019 A1	10/2021	Shua	
2018/0174062 A1	6/2018	Simo	2021/0336976 A1	10/2021	Shua	
2018/0181750 A1	6/2018	Lamothe-Brassard	2021/0377287 A1	12/2021	Shua	
2018/0191781 A1	7/2018	Palani et al.	2021/0406917 A1	12/2021	Erickson et al.	
2018/0211425 A1	7/2018	Winternitz et al.	2022/0004718 A1	1/2022	Quamar et al.	
2018/0219888 A1*	8/2018	Apostolopoulos	2022/0050840 A1	2/2022	Parravicini et al.	
		H04L 63/1425	2022/0058193 A1	2/2022	Smith et al.	
2018/0219897 A1	8/2018	Muddu et al.	2022/0067186 A1	3/2022	Thakur et al.	
2018/0227286 A1	8/2018	Ohsumi	2022/0086179 A1	3/2022	Levin et al.	
2018/0248901 A1	8/2018	Rieke	2022/0092481 A1	3/2022	Neithalath et al.	

(56)

References Cited**U.S. PATENT DOCUMENTS**

2022/0094614	A1	3/2022	Khurshid et al.
2022/0121741	A1	4/2022	Araujo et al.
2022/0124108	A1	4/2022	Gamble et al.
2022/0129803	A1	4/2022	Bikumala et al.
2022/0191226	A1	6/2022	Chan et al.
2022/0327119	A1	10/2022	Gasper et al.
2022/0342690	A1	10/2022	Shua
2022/0345480	A1	10/2022	Shua
2022/0345481	A1	10/2022	Shua
2022/0345483	A1	10/2022	Shua
2022/0350789	A1	11/2022	Yang et al.
2022/0350931	A1	11/2022	Shua
2022/0374800	A1	11/2022	Adinarayan et al.
2022/0376970	A1	11/2022	Chawathe et al.
2022/0382611	A1	12/2022	Kapish et al.
2022/0394082	A1	12/2022	Keren et al.
2022/0414072	A1	12/2022	Tandon et al.
2022/0414105	A1	12/2022	Umay et al.
2023/0025252	A1	1/2023	Erickson et al.
2023/0039566	A1	2/2023	Ghag et al.
2023/0052827	A1	2/2023	Araujo et al.
2023/0088960	A1	3/2023	Popelka et al.
2023/0096930	A1	3/2023	Dasdan
2023/0101773	A1	3/2023	Katahanas et al.
2023/0138371	A1	5/2023	Bandukwala et al.
2023/0244523	A1	8/2023	Gorantla et al.
2023/0251960	A1	8/2023	Sharma et al.
2023/0275909	A1	8/2023	Shivamoggi et al.
2023/0291755	A1	9/2023	Siebel et al.

FOREIGN PATENT DOCUMENTS

KR	20100133713	A	12/2010
WO	0227528	A1	4/2002

OTHER PUBLICATIONS

Mahajan et al., "A cloud based deployment framework for load balancing policies," 2014 Seventh International Conference on Contemporary Computing (IC3), Noida, India, 2014, pp. 565-570, doi: 10.1109/IC3.2014.6897235. (Year: 2014).*

Sakr et al., "A Survey of Large Scale Data Management Approaches in Cloud Environments," in IEEE Communications Surveys & Tutorials, vol. 13, No. 3, pp. 311-336, Third Quarter 2011, doi: 10.1109/SURV.2011.032211.00087. (Year: 2011).*

Tovarnák et al., "Continuous queries over distributed streams of heterogeneous monitoring data in cloud datacenters," 2014 9th International Conference on Software Engineering and Applications (ICSOFT-EA), Vienna, Austria, 2014, pp. 470-481. (Year: 2014).*

Oudejans et al., "QOMPLIANCE: Declarative Data-Centric Policy Compliance," 2023 IEEE Asia-Pacific Conference on Computer Science and Data Engineering (CSDE), Nadi, Fiji, 2023, pp. 01-08, doi: 10.1109/CSDE59766.2023.10487688. (Year: 2023).*

Anzures-García et al., "Policy-based Group Organizational Structure Management using an Ontological Approach," 2008 Third International Conference on Availability, Reliability and Security, Barcelona, Spain, 2008, pp. 807-812, doi: 10.1109/ARES.2008.186. (Year: 2008).*

Ghosh et al., "Verifying conformance of security implementation with organizational access policies in community cloud a formal approach," 2014 11th International Conference on Security and Cryptography (SECRYPT), Vienna, Austria, 2014, pp. 1-8. (Year: 2014).*

Paul et al., "Amazon Web Services Cloud Compliance Automation with Open Policy Agent," 2024 International Conference on Expert Clouds and Applications (ICOECA), Bengaluru, India, 2024, pp. 313-317, doi: 10.1109/ICOECA623512024.00063. (Year: 2024).*

Maurya et al., "Towards Data Gravity and Compliance Aware Distributed Deep Learning on Hybrid Clouds," 2022 IEEE 29th International Conference on High Performance Computing, Data

and Analytics Workshop (HiPCW), Bengaluru, India, 2022, pp. 53-58, doi: 10.1109/HiPCW57629.2022.00012. (Year: 2022).*

Singh et al., "Data Flow Management and Compliance in Cloud Computing," in IEEE Cloud Computing, vol. 2, No. 4, pp. 24-32, Jul.-Aug. 2015, doi: 10.1109/MCC.2015.69. (Year: 2015).*

Amidon, et al., "Program Fracture and Recombination for Efficient Automatic Code Reuse", 2015 IEEE High Performance Extreme Computing Conference (HPEC), (2015), pp. 1-6, doi: 10.1109/HPEC.2015.7396314.

Long, et al., Automatic Input Rectification, 2012 34th International Conference on Software Engineering (ICSE), (2012), pp. 80-90, doi: 10.1109/ICSE.2012.6227204.

Perkins, et al., "Automatically Patching Errors in Deployed Software", SOSP '09: Proceedings of the ACM SIGOPS 22nd symposium on Operating systems principles, Oct. 2009, p. 87-102, <https://doi.org/10.1145/1629575.1629585>.

Rinard, "Living in the Comfort Zone", OOPSLA'07, Oct. 21-25, 2007, Montreal, Quebec, Canada, pp. 611-622.

Rinard, "Manipulating Program Functionality to Eliminate Security Vulnerabilities", Moving Target Defense. Springer, New York, NY, (2011). pp. 109-115.

Samuel, et al., Let's Parse to Prevent Pwnage Invited Position Paper, LEET'12: Proceedings of the 5th USENIX conference on Large-Scale Exploits and Emergent Threats, Apr. 2012, pp. 3-6.

Shen, et al., "Active Learning for Inference and Regeneration of Applications that Access Databases", ACM Trans. Program. Lang. Syst. 42, 4, Article 18 (Jan. 2021), 119 pages, <https://doi.org/10.1145/3430952>.

Vasilakis, et al., "Supply-Chain Vulnerability Elimination via Active Learning and Regeneration", Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security (CCS '21), Nov. 15-19, 2021, Virtual Event, Republic of Korea. ACM, New York, NY, USA, 16 pages, <https://doi.org/10.1145/3460120.3484736>.

Ai-Yaseen, et al., "Real-time intrusion detection system using multi-agent system", IAENG International Journal of Computer Science 43, No. 1 (2016): 80-90.

Akoglu, et al., "Graph-based Anomaly Detection and Description: A Survey", Apr. 28, 2014.

Ammar, et al., "Query Optimization Techniques in Graph Databases", International Journal of Database Management Systems (IIDMS), vol. 8, No. 4, Aug. 2016, pp. 1-14 (Year: 2016).

Balasubramanian, et al., "An architecture for intrusion detection using autonomous agents", In Proceedings 14th annual computer security applications conference {Cal. No. 98EX217}, pp. 13-24. IEEE, 1998.

Beutel, et al., "User Behavior Modeling with Large-Scale Graph Analysis", Computer Science Department, Carnegie Mellon University, May 2016.

Bugiel, et al., "Towards Taming Privilege-Escalation Attacks on Android", In NOSS (vol. 17, p. 19). 2012, February.

Chang, et al., "Reality bites-progressive querying and result visualization in logical and VR spaces", doi: 10.1109/NL.1994.363635, 1994, pp. 100-109 (year 1994).

Chesson, et al., "Communication and control in a cluster network", ACM '74: Proceedings of the 1974 annual ACM conference—vol. 2, Jan. 1974, pp. 509-514, <http://doi.org/10.1145/1408839> (Year 1974).

Crosbie, et al., "Defending a computer system using autonomous agents", Department of Computer Science Technical Reports, Paper 1200, (1995).

Hautamaki, et al., "Outlier detection using k-nearest neighbour graph", Proceedings of the 17th International Conference on Pattern Recognition, 2004. ICPR 2004. vol. 3. IEEE, 2004.

Hooper, et al., "Medusa: a simple tool for interaction graph analysis", Bioinformatics, vol. 21 No. 24, 2005, pp. 1432-4433. {Year: 2005}.

Koutra, et al., "Exploring and Making Sense of Large Graphs", Computer Science Department, Carnegie Mellon University, Aug. 2015.

Leopold, et al., "A visual query system for the specification and scientific analysis of continual queries", doi: 10.1109/HCC.2001.995260, 2001, pp. 203-211, (Year: 2001).

(56)

References Cited**OTHER PUBLICATIONS**

Liao, et al., "Visualizing graph dynamics and similarity for enterprise network security and management", Proceedings of the seventh international symposium on visualization for cyber security. ACM, 2010.

Mateescu, et al., "Join-Graph Propagation Algorithm", Journal of Artificial Intelligence Research 37, 2010, pp. 279-328. (Year: 2010).

Moriano, et al., "Insider Threat Event Detection in User-System Interactions", MIST '17: Proceedings of the 2017 International Workshop on Managing Insider Security Threats, Oct. 2017, pp. 1-12 <https://doi.org/10.1145/3139923.3139928> (Year: 2017).

Ranshous, et al., "Anomaly detection in dynamic networks: a survey", WIREs Comput Stat, May/Jun. 2015.

Tamassia, et al., "Graph drawing for security visualization", International Symposium on Graph Drawing. Springer, Berlin, Heidelberg, 2008.

Vaas, et al., "Detecting disguised processes using application-behavior profiling", In 2017 IEEE International Symposium on Technologies for Homeland Security (HST), pp. 1-6. IEEE, 2017.

Yu, et al., "Recommending Join Queries Based on Path Frequency", IEEE, doi: 10.1109/WISA.2015.52, 2015, pp. 21-26. (Year: 2015).

Lim et al., "A policy-based approach for assuring data integrity in DBMSs," 6th International Conference on Collaborative Computing: Networking, Applications and Worksharing. 2010.

Mironeanu et al., "Intrusion detection using alert prioritization and multiple minimum supports," 14th RoEduNet International Conference. 2015.

Tong et al. "Finding Needles in a Moving Haystack: Prioritizing Alerts with Adversarial Reinforcement Learning," arXIV. Jun. 20, 2019.

Alsubhi et al., "Alert prioritization in intrusion Detection Systems," NOMS 2008 IEEE Network Operations and Management Symposium.

Bry et al., "A High-Level Query Language for Events", 2006 IEEE Services Computing Workshops.

Lagraa et al., "What database do you choose for heterogeneous security log events analysis?", 2021 IFIP/IEEE International Symposium on Integrated Network Management (IM).

Mironeanu et al., "Intrusion detection using alert prioritization and multiple minimum supports," 2005 14th RoEduNet International Conference—Networking in Education and Research.

Schuster et al., "Control-Flow-Based Querying of Process Executions from Partially Ordered Event Data", Nov. 8, 2022.

Swinnen et al., "Querying event logs: Discovering non-events in event logs", 2010 IEEE International Conference on Intelligent Systems and Knowledge Engineering.

Yurcik et al., "UCLog+: A Security Data Management System for Correlating Alerts, Incidents, and Raw Data From Remote Logs," Jul. 25, 2006.

Office Action for U.S. Appl. No. 17/989,370 mailed Feb. 18, 2025, 14 pages.

* cited by examiner




Fig. 1A




Fig. 1B




Fig. 1C




Fig. 1D




Fig. 2A




Fig. 2B




Fig. 2C




Fig. 2D






Fig. 2E








Fig. 2H




Fig. 21




Fig. 2J




Fig. 2K




Fig. 2L




Fig. 2M




Fig. 2N




Fig. 20








Fig. 3A






Fig. 3F




Fig. 3G




Fig. 3H




Fig. 3I




Fig. 3J

MID	start_time	PID_hash	src_IP_addr	src_port	dst_IP_addr	dst_port	prot	dir
A	t1	A1	1.1.1.10	10000	2.2.2.20	22	TCP	Incoming
A	t2	A3	2.2.2.20	10001	2.2.2.21	22	TCP	Outgoing
B	t2	B1	2.2.2.20	10001	2.2.2.21	22	TCP	Incoming

Fig. 3K

src_MID	src_PID_hash	dst_MID	dst_PID_hash	dst_start_time	src_IP_addr	src_port	dst_IP_addr	dst_port
null	null	A	A1	t1	1.1.1.10	10000	2.2.2.20	22
A	A3	B	B1	t2	2.2.2.20	10001	2.2.2.21	22

Fig. 3L

MID	login_time	sshd_PID_hash
A	t1	A1
B	t2	B1

Fig. 3M

MID	sshd_PID_hash	login_time	login_username	src_IP_addr	src_port	dst_IP_addr	dst_port
A	A1	t1	X	1.1.1.10	10000	2.2.2.20	22
B	B1	t2	Y	2.2.2.20	10001	2.2.2.21	22

Fig. 3N

MID	start_time	PID_hash	exe_path	parent_PID_hash
A	t1	A1	/usr/sbin/sshd	A0
A	t1	A2	/bin/bash	A1
A	t2	A3	/usr/bin/ssh	A2
B	t2	B1	/usr/sbin/sshd	B0
B	t2	B2	/bin/bash	B1
B	t3	B3	/usr/bin/curl	B2

Fig. 30

MID	sshd_PID_hash	PID_hash
A	A1	A1
A	A1	A2
A	A1	A3
B	B1	B1
B	B1	B2
B	B1	B3

Fig. 3P

parent_MID	parent_sshd_PID_hash	child_MID	origin_sshd_PID_hash
A	A1	B	B1

Fig. 3Q

MID	sshd_PID_hash	parent_MID	parent_sshd_PID_hash	origin_MID	origin_sshd_PID_hash
A	A1	null	null	A	A1
B	B1	A	A1	A	A1

Fig. 3R




Fig. 3S




Fig. 4A




Fig. 4B




Fig. 4C




Fig. 4D





Fig. 4E



450






Fig. 4G


Fig. 4H



466



477






Fig. 4K




Fig. 4L




Fig. 5A




Fig. 5B




Fig. 5C




Fig. 5D

531

```
Global_AWS_CTA_LoginUnknownIP {  
  source {  
    CloudTrailRawEvents  
  }  
  
  filter {  
    EVENT_SOURCE = 'signin.amazonaws.com'  
    and EVENT:userIdentity.type::String = 'AWSService'  
    and EVENT:sourceIPAddress not in ('1.1.1.1', '2.2.2.2')  
    and ERROR_CODE is null  
  }  
  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

532

```
{  
  "policyType": "Violation",  
  "queryId": "Global_AWS_CTA_LoginUnknownIP",  
  "title": "Service Account Login from Unknown IP Address",  
  "enabled": true,  
  "description": "A service account logged in from an \  
    unexpected IP address",  
  "remediation": "Investigate any suspicious activity",  
  "severity": "high",  
  "limit": 100,  
  "evalFrequency": "Hourly",  
  "alertEnabled": true,  
  "alertProfile": "LW_CloudTrail_Alerts",  
  "policyUi": {  
    "domain": "AWS",  
    "subdomain": "Cloudtrail"  
  }  
  "evaluatorId": "Cloudtrail"  
}
```

Fig. 5E

533

```
LW_Custom_AWS_CTA_AuroraPasswordChange {  
  source {  
    CloudTrailRawEvents  
  }  
  filter {  
    EVENT_SOURCE = 'rds.amazonaws.com'  
    and EVENT_NAME = 'ModifyDBCluster'  
    and EVENT:requestParameters.masterUserPassword is not \  
      json null  
    and EVENT:requestParameters.applyImmediately = true  
    and ERROR_CODE is null  
  }  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

534

```
policies:  
- policyId: CustomPolicy1  
  title: Aurora Password Change  
  enabled: false  
  policyType: Violation  
  alertEnabled: false  
  alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert  
  evalFrequency: Hourly  
  queryId: LW_Custom_AWS_CTA_AuroraPasswordChange  
  limit: 100  
  severity: high  
  description: Password for an Aurora RDS cluster was changed  
  remediation: |-  
    Check that the password change was expected and ensure \  
    only specified users can modify the RDS cluster.  
  policyUi:  
    domain: AWS  
    subdomain: CloudTrail  
  evaluatorId: Cloudtrail
```

Fig. 5F

535

```
LW_Custom_AWS_CTA_RDSBackupDisabled {  
  source {  
    CloudTrailRawEvents  
  }  
  filter {  
    EVENT_SOURCE = 'rds.amazonaws.com'  
    and EVENT_NAME = 'ModifyDBInstance'  
    and EVENT:requestParameters.backupRetentionPeriod = 0  
    and ERROR_CODE is null  
  }  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

536

```
policies:  
- policyId: CustomPolicy2  
  title: RDS Backup Disabled  
  enabled: false  
  policyType: Violation  
  alertEnabled: false  
  alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert  
  evalFrequency: Hourly  
  queryId: LW_Custom_AWS_CTA_RDSBackupDisabled  
  limit: 100  
  severity: high  
  description: Backups have been disabled for an RDS database\  
    instance  
  remediation: |-  
    Ensure only specified users can modify DB instances. \  
    To reenale, navigate to the database, select Modify, \  
    set Backup Retention Period greater than 0.  
  policyUi:  
    domain: AWS  
    subdomain: CloudTrail  
  evaluatorId: Cloudtrail
```

Fig. 5G

537

```
LW_Custom_AWS_CTA_RootAccountMFADisabled {
  source {
    CloudTrailRawEvents
  }
  filter {
    EVENT_SOURCE = 'iam.amazonaws.com'
    and ((EVENT_NAME = 'DeactivateMFADevice'
    and EVENT:requestParameters.userName = 'AWS ROOT USER')
    or EVENT_NAME = 'DeleteVirtualMFADevice')
    and EVENT:requestParameters.serialNumber like \
      '%mfa/root-account-mfa-device'
    and ERROR_CODE is null
  }
  return distinct {
    INSERT_ID
    INSERT_TIME,
    EVENT_TIME
    EVENT
  }
}
```

538

```
policies:
- policyId: CustomPolicy3
  title: AWS Root Account MFA Disabled
  enabled: false
  policyType: Violation
  alertEnabled: false
  alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert
  evalFrequency: Hourly
  queryId: LW_Custom_AWS_CTA_RootAccountMFADisabled
  limit: 100
  severity: critical
  description: AWS root account MFA device was deactivated or\
    deleted
  remediation: |-
    Enable MFA for the root account: \
    1. Log in to your account using the root credentials. \
    2. Choose the account name near the top-right corner... \
    /* ADDITIONAL STEPS */
  policyUi:
    domain: AWS
    subdomain: CloudTrail
  evaluatorId: Cloudtrail
```

Fig. 5H

539

```
LW_Custom_AWS_CTA_KMSChange {  
  source {  
    CloudTrailRawEvents  
  }  
  filter {  
    EVENT_SOURCE = 'kms.amazonaws.com'  
    and ERROR_CODE is null  
  }  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

540

```
policies:  
  - policyId: CustomPolicy4  
    title: KMS Change  
    enabled: false  
    policyType: Violation  
    alertEnabled: false  
    alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert  
    evalFrequency: Hourly  
    queryId: LW_Custom_AWS_CTA_KMSChange  
    limit: 100  
    severity: high  
    description: A KMS change occurred  
    remediation: |-  
      Check the KMS change and ensure only specified users \  
      can make KMS changes  
    policyUi:  
      domain: AWS  
      subdomain: CloudTrail  
    evaluatorId: Cloudtrail
```

Fig. 51

541

```
LW_Custom_AWS_CTA_FailedLogin {  
  source {  
    CloudTrailRawEvents  
  }  
  filter {  
    EVENT_SOURCE = 'signin.amazonaws.com'  
    and EVENT_NAME in ('ConsoleLogin')  
    and EVENT:responseElements.ConsoleLogin::String = \  
      'Failure'  
    and ERROR_CODE is null  
  }  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

542

```
policies:  
- policyId: CustomPolicy5  
  title: Failed Login  
  enabled: false  
  policyType: Violation  
  alertEnabled: false  
  alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert  
  evalFrequency: Hourly  
  queryId: LW_Custom_AWS_CTA_FailedLogin  
  limit: 100  
  severity: medium  
  description: A failed console login attempt was detected  
  remediation: |-  
    Investigate any suspicious activity  
  policyUi:  
    domain: AWS  
    subdomain: CloudTrail  
  evaluatorId: Cloudtrail
```

Fig. 5J

543

```
LW_Custom_AWS_CTA_UnencryptedVolumeCreated {  
  source {  
    CloudTrailRawEvents  
  }  
  filter {  
    EVENT_SOURCE = 'ec2.amazonaws.com'  
    and EVENT_NAME = 'CreateVolume'  
    and EVENT:responseElements.encrypted = false  
    and ERROR_CODE is null  
  }  
  return distinct {  
    INSERT_ID  
    INSERT_TIME,  
    EVENT_TIME  
    EVENT  
  }  
}
```

544

```
policies:  
  - policyId: CustomPolicy6  
    title: Unencrypted Volume Created  
    enabled: false  
    policyType: Violation  
    alertEnabled: false  
    alertProfile: LW_CloudTrail_Alerts.CloudTrailDefaultAlert  
    evalFrequency: Hourly  
    queryId: LW_Custom_AWS_CTA_UnencryptedVolumeCreated  
    limit: 100  
    severity: high  
    description: It is a best practice to encrypt EBS volumes  
    remediation: |-  
      It is not possible to encrypt an unencrypted EBS volume.\n      We suggest creating a snapshot, then creating an \n      encrypted volume with the snapshot.  
    policyUi:  
      domain: AWS  
      subdomain: CloudTrail  
    evaluatorId: Cloudtrail
```

Fig. 5K




Fig. 6A




Fig. 6B




Fig. 6C

631

```

LW_Profile: {
  fields: {
    /* define fields such as Number, String, Timestamp, IP_ADDR, PORT, PROTOCOL,
      AWS_ACCOUNT_ID, etc. */
  }
  entityKeys: {
    /* define entity keys such as ProcessKey, MachineKey, ConnectionKey, etc. */
  }
}

```

632

```

LW_AlertProfile: {
  extends: LW_Profile
  fields: {
    /* define new fields and extend fields inherited from LW_Profile */
  }
}

```

633

```

LW_AWS_AlertProfile: {
  extends: LW_AlertProfile,
  fields: {
    /* define new fields and extend fields inherited from LW_AlertProfile and/or
LW_Profile */
  }
  entityKeys: [
    { /* define entity keys such as aws_resource_key */ },
  ],
  descriptionKeys: [
    { /* define description key 1 (e.g., src_accountid_desc_key) */ },
    { /* define description key 2 (e.g., src_recipient_alias_desc_key) */ },
    { /* define description key 3 (e.g., src_recid_desc_key) */ },
    { /* define other description keys... (e.g., src_title_desc_key, etc.) */ },
  ]
  alerts: [
    {
      name: "NewViolation",
      eventType: "New Violations",
      eventName: "NewViolations"
      description: "AWS Account /* reference description keys */"
      subject: "/* reference description keys */ violations detected"
    },
    { /* define one or more additional alert definitions */ }
  ]
}

```

Fig. 6D

634

```
policies:
- policyId: LQL_S3_Policy_1
  title: User Defined LQL Policy 1 for S3
  enabled: true
  policyType: Violation
  alertEnabled: true
  alertProfile: LW_AWS_AlertProfile
  evalFrequency: Hourly
  evalSummary: LQL_S3_SUMMARY
  queryId: LQL_S3_Q1
  limit: 1000
  severity: medium
  description: /* Policy description */
  remediation: /* Policy remediation */
```

635

```
policies:
- policyId: LQL_S3_Policy_2
  title: User Defined LQL Policy 2 for S3
  enabled: true
  policyType: Violation
  alertEnabled: true
  alertProfile:
    name: LW_AWS_AlertProfile
    alerts.ComplianceChanged.enabled:false
    fields.POLICY_TITLE.value:LQL_S3_Policy2_XXX
    fieldTypeMap: /* define field type info for returned JSON fields */
    CONFIG.x.y.z: IP_ADDR
  evalFrequency: Hourly
  evalSummary: LQL_S3_SUMMARY
  queryId: LQL_S3_Q1
  limit: 1000
  severity: medium
  description: /* Policy description */
  remediation: /* Policy remediation */
```

636

Fig. 6E




Fig. 7A




Fig. 7B




Fig. 7C

COMMUNITY-BASED GENERATION OF POLICIES FOR A DATA PLATFORM

RELATED APPLICATIONS

The present application is a continuation-in-part application of U.S. patent application Ser. No. 17/510,179 filed on Oct. 25, 2021, which is a continuation of U.S. patent application Ser. No. 16/786,822 filed Feb. 10, 2020, now U.S. Pat. No. 11,157,502, which is a continuation of U.S. patent application Ser. No. 16/134,806 filed Sep. 18, 2018, now U.S. Pat. No. 10,614,071, which claims priority to U.S. Provisional Patent Application No. 62/590,986 filed Nov. 27, 2017 and U.S. Provisional Patent Application No. 62/650,971 filed Mar. 30, 2018. The present application also claims priority to U.S. Provisional Patent Application No. 63/224,373, filed Jul. 21, 2021; U.S. Provisional Patent Application No. 63/228,980, filed Jul. 21, 2021; U.S. Provisional Patent Application No. 63/239,275, filed Aug. 31, 2021; and U.S. Provisional Patent Application No. 63/248,915, filed Sep. 27, 2021. The contents of these applications are hereby incorporated by reference in their entirety.

BRIEF DESCRIPTION OF THE DRAWINGS

The accompanying drawings illustrate various embodiments and are a part of the specification. The illustrated embodiments are merely examples and do not limit the scope of the disclosure. Throughout the drawings, identical or similar reference numbers designate identical or similar elements.

FIG. 1A shows an illustrative configuration in which a data platform is configured to perform various operations with respect to a cloud environment that includes a plurality of compute assets.

FIG. 1B shows an illustrative implementation of the configuration of FIG. 1A.

FIG. 1C illustrates an example computing device.

FIG. 1D illustrates an example of an environment in which activities that occur within datacenters are modeled.

FIG. 2A illustrates an example of a process, used by an agent, to collect and report information about a client.

FIG. 2B illustrates a 5-tuple of data collected by an agent, physically and logically.

FIG. 2C illustrates a portion of a polygraph.

FIG. 2D illustrates a portion of a polygraph.

FIG. 2E illustrates an example of a communication polygraph.

FIG. 2F illustrates an example of a polygraph.

FIG. 2G illustrates an example of a polygraph as rendered in an interface.

FIG. 2H illustrates an example of a portion of a polygraph as rendered in an interface.

FIG. 2I illustrates an example of a portion of a polygraph as rendered in an interface.

FIG. 2J illustrates an example of a portion of a polygraph as rendered in an interface.

FIG. 2K illustrates an example of a portion of a polygraph as rendered in an interface.

FIG. 2L illustrates an example of an insider behavior graph as rendered in an interface.

FIG. 2M illustrates an example of a privilege change graph as rendered in an interface.

FIG. 2N illustrates an example of a user login graph as rendered in an interface.

FIG. 2O illustrates an example of a machine server graph as rendered in an interface.

FIG. 3A illustrates an example of a process for detecting anomalies in a network environment.

FIG. 3B depicts a set of example processes communicating with other processes.

FIG. 3C depicts a set of example processes communicating with other processes.

FIG. 3D depicts a set of example processes communicating with other processes.

FIG. 3E depicts two pairs of clusters.

FIG. 3F is a representation of a user logging into a first machine, then into a second machine from the first machine, and then making an external connection.

FIG. 3G is an alternate representation of actions occurring in FIG. 3F.

FIG. 3H illustrates an example of a process for performing extended user tracking.

FIG. 3I is a representation of a user logging into a first machine, then into a second machine from the first machine, and then making an external connection.

FIG. 3J illustrates an example of a process for performing extended user tracking.

FIG. 3K illustrates example records.

FIG. 3L illustrates example output from performing an ssh connection match.

FIG. 3M illustrates example records.

FIG. 3N illustrates example records.

FIG. 3O illustrates example records.

FIG. 3P illustrates example records.

FIG. 3Q illustrates an adjacency relationship between two login sessions.

FIG. 3R illustrates example records.

FIG. 3S illustrates an example of a process for detecting anomalies.

FIG. 4A illustrates a representation of an embodiment of an insider behavior graph.

FIG. 4B illustrates an embodiment of a portion of an insider behavior graph.

FIG. 4C illustrates an embodiment of a portion of an insider behavior graph.

FIG. 4D illustrates an embodiment of a portion of an insider behavior graph.

FIG. 4E illustrates a representation of an embodiment of a user login graph.

FIG. 4F illustrates an example of a privilege change graph.

FIG. 4G illustrates an example of a privilege change graph.

FIG. 4H illustrates an example of a user interacting with a portion of an interface.

FIG. 4I illustrates an example of a dossier for an event.

FIG. 4J illustrates an example of a dossier for a domain.

FIG. 4K depicts an example of an Entity Join graph by FilterKey and FilterKey Group (implicit join).

FIG. 4L illustrates an example of a process for dynamically generating and executing a query.

FIG. 5A illustrates an example of a process for implementing custom policies and queries for a data platform.

FIG. 5B illustrates an example policy that invokes a query and an alert profile to provide an annotated query result to a user.

FIG. 5C illustrates examples of different types of policies invoking different types of queries.

FIG. 5D illustrates an example configuration in which a data platform may query different types of target datasets at different stages of a process involving ingesting and processing data from a cloud environment.

3

FIG. 5E illustrates an example of a query and a custom policy that invokes the query.

FIG. 5F illustrates an example of a query and a custom policy that invokes the query.

FIG. 5G illustrates an example of a query and a custom policy that invokes the query.

FIG. 5H illustrates an example of a query and a custom policy that invokes the query.

FIG. 5I illustrates an example of a query and a custom policy that invokes the query.

FIG. 5J illustrates an example of a query and a custom policy that invokes the query.

FIG. 5K illustrates an example of a query and a custom policy that invokes the query.

FIG. 6A illustrates an example of a process for implementing modular alert profiles for custom policies and queries.

FIG. 6B illustrates an example of how different custom policies may invoke a same alert profile.

FIG. 6C illustrates an example architecture by way of which custom policies invoking a same alert profile generate alerts for presentation to a user.

FIG. 6D illustrates example alert profiles and how they may inherit attributes from antecedent alert profiles.

FIG. 6E illustrates example custom policies that invoke a particular alert profile with and without overriding an attribute of the alert profile.

FIG. 7A illustrates an example of a process for implementing community-based generation of policies for a data platform.

FIG. 7B illustrates an example configuration in which community-based generation of custom policies is implemented.

FIG. 7C illustrates example aspects of how a community policy may be generated.

DETAILED DESCRIPTION

Various illustrative embodiments are described herein with reference to the accompanying drawings. It will, however, be evident that various modifications and changes may be made thereto, and additional embodiments may be implemented, without departing from the scope of the invention as set forth in the claims. For example, certain features of one embodiment described herein may be combined with or substituted for features of another embodiment described herein. The description and drawings are accordingly to be regarded in an illustrative rather than a restrictive sense.

FIG. 1A shows an illustrative configuration 10 in which a data platform 12 is configured to perform various operations with respect to a cloud environment 14 that includes a plurality of compute assets 16-1 through 16-N (collectively “compute assets 16”). For example, data platform 12 may include data ingestion resources 18 configured to ingest data from cloud environment 14 into data platform 12, data processing resources 20 configured to perform data processing operations with respect to the data, and user interface resources 22 configured to provide one or more external users and/or compute resources (e.g., computing device 24) with access to an output of data processing resources 20. Each of these resources are described in detail herein.

Cloud environment 14 may include any suitable network-based computing environment as may serve a particular application. For example, cloud environment 14 may be implemented by one or more compute resources provided and/or otherwise managed by one or more cloud service providers, such as Amazon Web Services (AWS), Google

4

Cloud Platform (GCP), Microsoft Azure, and/or any other cloud service provider configured to provide public and/or private access to network-based compute resources.

Compute assets 16 may include, but are not limited to, containers (e.g., container images, deployed and executing container instances, etc.), virtual machines, workloads, applications, processes, physical machines, compute nodes, clusters of compute nodes, software runtime environments (e.g., container runtime environments), and/or any other virtual and/or physical compute resource that may reside in and/or be executed by one or more computer resources in cloud environment 14. In some examples, one or more compute assets 16 may reside in one or more datacenters.

A compute asset 16 may be associated with (e.g., owned, deployed, or managed by) a particular entity, such as a customer or client of cloud environment 14 and/or data platform 12. Accordingly, for purposes of the discussion herein, cloud environment 14 may be used by one or more entities.

Data platform 12 may be configured to perform one or more data security monitoring and/or remediation services, compliance monitoring services, anomaly detection services, DevOps services, compute asset management services, and/or any other type of data analytics service as may serve a particular implementation. Data platform 12 may be managed or otherwise associated with any suitable data platform provider, such as a provider of any of the data analytics services described herein. The various resources included in data platform 12 may reside in the cloud and/or be located on-premises and be implemented by any suitable combination of physical and/or virtual compute resources, such as one or more computing devices, microservices, applications, etc.

Data ingestion resources 18 may be configured to ingest data from cloud environment 14 into data platform 12. This may be performed in various ways, some of which are described in detail herein. For example, as illustrated by arrow 26, data ingestion resources 18 may be configured to receive the data from one or more agents deployed within cloud environment 14, utilize an event streaming platform (e.g., Kafka) to obtain the data, and/or pull data (e.g., configuration data) from cloud environment 14. In some examples, data ingestion resources 18 may obtain the data using one or more agentless configurations.

The data ingested by data ingestion resources 18 from cloud environment 14 may include any type of data as may serve a particular implementation. For example, the data may include data representative of configuration information associated with compute assets 16, information about one or more processes running on compute assets 16, network activity information, information about events (creation events, modification events, communication events, user-initiated events, etc.) that occur with respect to compute assets 16, etc. In some examples, the data may or may not include actual customer data processed or otherwise generated by compute assets 16.

As illustrated by arrow 28, data ingestion resources 18 may be configured to load the data ingested from cloud environment 14 into a data store 30. Data store 30 is illustrated in FIG. 1A as being separate from and communicatively coupled to data platform 12. However, in some alternative embodiments, data store 30 is included within data platform 12.

Data store 30 may be implemented by any suitable data warehouse, data lake, data mart, and/or other type of database structure as may serve a particular implementation. Such data stores may be proprietary or may be embodied as

vendor provided products or services such as, for example, Snowflake, Google BigQuery, Druid, Amazon Redshift, IBM Db2, Dremio, Databricks Lakehouse Platform, Cloudera, Azure Synapse Analytics, and others.

Although the examples described herein largely relate to embodiments where data is collected from agents and ultimately stored in a data store such as those provided by Snowflake, in other embodiments data that is collected from agents and other sources may be stored in different ways. For example, data that is collected from agents and other sources may be stored in a data warehouse, data lake, data mart, and/or any other data store.

A data warehouse may be embodied as an analytic database (e.g., a relational database) that is created from two or more data sources. Such a data warehouse may be leveraged to store historical data, often on the scale of petabytes. Data warehouses may have compute and memory resources for running complicated queries and generating reports. Data warehouses may be the data sources for business intelligence ('BI') systems, machine learning applications, and/or other applications. By leveraging a data warehouse, data that has been copied into the data warehouse may be indexed for good analytic query performance, without affecting the write performance of a database (e.g., an Online Transaction Processing ('OLTP') database). Data warehouses also enable joining data from multiple sources for analysis. For example, a sales OLTP application probably has no need to know about the weather at various sales locations, but sales predictions could take advantage of that data. By adding historical weather data to a data warehouse, it would be possible to factor it into models of historical sales data.

Data lakes, which store files of data in their native format, may be considered as "schema on read" resources. As such, any application that reads data from the lake may impose its own types and relationships on the data. Data warehouses, on the other hand, are "schema on write," meaning that data types, indexes, and relationships are imposed on the data as it is stored in an enterprise data warehouse (EDW). "Schema on read" resources may be beneficial for data that may be used in several contexts and poses little risk of losing data. "Schema on write" resources may be beneficial for data that has a specific purpose, and good for data that must relate properly to data from other sources. Such data stores may include data that is encrypted using homomorphic encryption, data encrypted using privacy-preserving encryption, smart contracts, non-fungible tokens, decentralized finance, and other techniques.

Data marts may contain data oriented towards a specific business line whereas data warehouses contain enterprise-wide data. Data marts may be dependent on a data warehouse, independent of the data warehouse (e.g., drawn from an operational database or external source), or a hybrid of the two. In embodiments described herein, different types of data stores (including combinations thereof) may be leveraged.

Data processing resources 20 may be configured to perform various data processing operations with respect to data ingested by data ingestion resources 18, including data ingested and stored in data store 30. For example, data processing resources 20 may be configured to perform one or more data security monitoring and/or remediation operations, compliance monitoring operations, anomaly detection operations, DevOps operations, compute asset management operations, and/or any other type of data analytics operation as may serve a particular implementation. Various examples of operations performed by data processing resources 20 are described herein.

As illustrated by arrow 32, data processing resources 20 may be configured to access data in data store 30 to perform the various operations described herein. In some examples, this may include performing one or more queries with respect to the data stored in data store 30. Such queries may be generated using any suitable query language.

In some examples, the queries provided by data processing resources 20 may be configured to direct data store 30 to perform one or more data analytics operations with respect to the data stored within data store 30. These data analytics operations may be with respect to data specific to a particular entity (e.g., data residing in one or more silos within data store 30 that are associated with a particular customer) and/or data associated with multiple entities. For example, data processing resources 20 may be configured to analyze data associated with a first entity and use the results of the analysis to perform one or more operations with respect to a second entity.

One or more operations performed by data processing resources 20 may be performed periodically according to a predetermined schedule. For example, one or more operations may be performed by data processing resources 20 every hour or any other suitable time interval. Additionally or alternatively, one or more operations performed by data processing resources 20 may be performed in substantially real-time (or near real-time) as data is ingested into data platform 12. In this manner, the results of such operations (e.g., one or more detected anomalies in the data) may be provided to one or more external entities (e.g., computing device 24 and/or one or more users) in substantially real-time and/or in near real-time.

User interface resources 22 may be configured to perform one or more user interface operations, examples of which are described herein. For example, user interface resources 22 may be configured to present one or more results of the data processing performed by data processing resources 20 to one or more external entities (e.g., computing device 24 and/or one or more users), as illustrated by arrow 34. As illustrated by arrow 36, user interface resources 22 may access data in data store 30 to perform the one or more user interface operations.

FIG. 1B illustrates an implementation of configuration 10 in which an agent 38 (e.g., agent 38-1 through agent 38-N) is installed on each of compute assets 16. As used herein, an agent may include a self-contained binary and/or other type of code or application that can be run on any appropriate platforms, including within containers and/or other virtual compute assets. Agents 38 may monitor the nodes on which they execute for a variety of different activities, including but not limited to, connection, process, user, machine, and file activities. In some examples, agents 38 can be executed in user space, and can use a variety of kernel modules (e.g., auditd, iptables, netfilter, pcap, etc.) to collect data. Agents can be implemented in any appropriate programming language, such as C or Golang, using applicable kernel APIs.

Agents 38 may be deployed in any suitable manner. For example, an agent 38 may be deployed as a containerized application or as part of a containerized application. As described herein, agents 38 may selectively report information to data platform 12 in varying amounts of detail and/or with variable frequency.

Also shown in FIG. 1B is a load balancer 40 configured to perform one or more load balancing operations with respect to data ingestion operations performed by data ingestion resources 18 and/or user interface operations performed by user interface resources 22. Load balancer 40 is shown to be included in data platform 12. However, load

balancer **40** may alternatively be located external to data platform **12**. Load balancer **40** may be implemented by any suitable microservice, application, and/or other computing resources. In some alternative examples, data platform **12** may not utilize a load balancer such as load balancer **40**.

Also shown in FIG. **1B** is long term storage **42** with which data ingestion resources **18** may interface, as illustrated by arrow **44**. Long term storage **42** may be implemented by any suitable type of storage resources, such as cloud-based storage (e.g., AWS S3, etc.) and/or on-premises storage and may be used by data ingestion resources **18** as part of the data ingestion process. Examples of this are described herein. In some examples, data platform **12** may not utilize long term storage **42**.

The embodiments described herein can be implemented in numerous ways, including as a process; an apparatus; a system; a composition of matter; a computer program product embodied on a computer readable storage medium; and/or a processor, such as a processor configured to execute instructions stored on and/or provided by a memory coupled to the processor. In this specification, these implementations, or any other form that the invention may take, may be referred to as techniques. In general, the order of the steps of disclosed processes may be altered within the scope of the principles described herein. Unless stated otherwise, a component such as a processor or a memory described as being configured to perform a task may be implemented as a general component that is temporarily configured to perform the task at a given time or a specific component that is manufactured to perform the task. As used herein, the term ‘processor’ refers to one or more devices, circuits, and/or processing cores configured to process data, such as computer program instructions.

In some examples, a non-transitory computer-readable medium storing computer-readable instructions may be provided in accordance with the principles described herein. The instructions, when executed by a processor of a computing device, may direct the processor and/or computing device to perform one or more operations, including one or more of the operations described herein. Such instructions may be stored and/or transmitted using any of a variety of known computer-readable media.

A non-transitory computer-readable medium as referred to herein may include any non-transitory storage medium that participates in providing data (e.g., instructions) that may be read and/or executed by a computing device (e.g., by a processor of a computing device). For example, a non-transitory computer-readable medium may include, but is not limited to, any combination of non-volatile storage media and/or volatile storage media. Exemplary non-volatile storage media include, but are not limited to, read-only memory, flash memory, a solid-state drive, a magnetic storage device (e.g. a hard disk, a floppy disk, magnetic tape, etc.), ferroelectric random-access memory (“RAM”), and an optical disc (e.g., a compact disc, a digital video disc, a Blu-ray disc, etc.). Exemplary volatile storage media include, but are not limited to, RAM (e.g., dynamic RAM).

FIG. **1C** illustrates an example computing device **50** that may be specifically configured to perform one or more of the processes described herein. Any of the systems, microservices, computing devices, and/or other components described herein may be implemented by computing device **50**.

As shown in FIG. **1C**, computing device **50** may include a communication interface **52**, a processor **54**, a storage device **56**, and an input/output (“I/O”) module **58** communicatively connected one to another via a communication

infrastructure **60**. While an exemplary computing device **50** is shown in FIG. **1C**, the components illustrated in FIG. **1C** are not intended to be limiting. Additional or alternative components may be used in other embodiments. Components of computing device **50** shown in FIG. **1C** will now be described in additional detail.

Communication interface **52** may be configured to communicate with one or more computing devices. Examples of communication interface **52** include, without limitation, a wired network interface (such as a network interface card), a wireless network interface (such as a wireless network interface card), a modem, an audio/video connection, and any other suitable interface.

Processor **54** generally represents any type or form of processing unit capable of processing data and/or interpreting, executing, and/or directing execution of one or more of the instructions, processes, and/or operations described herein. Processor **54** may perform operations by executing computer-executable instructions **62** (e.g., an application, software, code, and/or other executable data instance) stored in storage device **56**.

Storage device **56** may include one or more data storage media, devices, or configurations and may employ any type, form, and combination of data storage media and/or device. For example, storage device **56** may include, but is not limited to, any combination of the non-volatile media and/or volatile media described herein. Electronic data, including data described herein, may be temporarily and/or permanently stored in storage device **56**. For example, data representative of computer-executable instructions **62** configured to direct processor **54** to perform any of the operations described herein may be stored within storage device **56**. In some examples, data may be arranged in one or more databases residing within storage device **56**.

I/O module **58** may include one or more I/O modules configured to receive user input and provide user output. I/O module **58** may include any hardware, firmware, software, or combination thereof supportive of input and output capabilities. For example, I/O module **58** may include hardware and/or software for capturing user input, including, but not limited to, a keyboard or keypad, a touchscreen component (e.g., touchscreen display), a receiver (e.g., an RF or infrared receiver), motion sensors, and/or one or more input buttons.

I/O module **58** may include one or more devices for presenting output to a user, including, but not limited to, a graphics engine, a display (e.g., a display screen), one or more output drivers (e.g., display drivers), one or more audio speakers, and one or more audio drivers. In certain embodiments, I/O module **58** is configured to provide graphical data to a display for presentation to a user. The graphical data may be representative of one or more graphical user interfaces and/or any other graphical content as may serve a particular implementation.

FIG. **1D** illustrates an example implementation **100** of configuration **10**. As such, one or more components shown in FIG. **1D** may implement one or more components shown in FIG. **1A** and/or FIG. **1B**. In particular, implementation **100** illustrates an environment in which activities that occur within datacenters are modeled using data platform **12**. Using techniques described herein, a baseline of datacenter activity can be modeled, and deviations from that baseline can be identified as anomalous. Anomaly detection can be beneficial in a security context, a compliance context, an asset management context, a DevOps context, and/or any other data analytics context as may serve a particular implementation.

Two example datacenters (**104** and **106**) are shown in FIG. 1D, and are associated with (e.g., belong to) entities named entity A and entity B, respectively. A datacenter may include dedicated equipment (e.g., owned and operated by entity A, or owned/leased by entity A and operated exclusively on entity A's behalf by a third party). A datacenter can also include cloud-based resources, such as infrastructure as a service (IaaS), platform as a service (PaaS), and/or software as a service (SaaS) elements. The techniques described herein can be used in conjunction with multiple types of datacenters, including ones wholly using dedicated equipment, ones that are entirely cloud-based, and ones that use a mixture of both dedicated equipment and cloud-based resources.

Both datacenter **104** and datacenter **106** include a plurality of nodes, depicted collectively as set of nodes **108** and set of nodes **110**, respectively, in FIG. 1D. These nodes may implement compute assets **16**. Installed on each of the nodes are in-server/in-virtual-machine (VM)/embedded-in-IoT device agents (e.g., agent **112**), which are configured to collect data and report it to data platform **12** for analysis. As described herein, agents may be small, self-contained binaries that can be run on any appropriate platforms, including virtualized ones (and, as applicable, within containers). Agents may monitor the nodes on which they execute for a variety of different activities, including: connection, process, user, machine, and file activities. Agents can be executed in user space, and can use a variety of kernel modules (e.g., auditd, iptables, netfilter, pcap, etc.) to collect data. Agents can be implemented in any appropriate programming language, such as C or Golang, using applicable kernel APIs.

As described herein, agents can selectively report information to data platform **12** in varying amounts of detail and/or with variable frequency. As is also described herein, the data collected by agents may be used by data platform **12** to create polygraphs, which are graphs of logical entities, connected by behaviors. In some embodiments, agents report information directly to data platform **12**. In other embodiments, at least some agents provide information to a data aggregator, such as data aggregator **114**, which in turn provides information to data platform **12**. The functionality of a data aggregator can be implemented as a separate binary or other application (distinct from an agent binary), and can also be implemented by having an agent execute in an "aggregator mode" in which the designated aggregator node acts as a Layer 7 proxy for other agents that do not have access to data platform **12**. Further, a chain of multiple aggregators can be used, if applicable (e.g., with agent **112** providing data to data aggregator **114**, which in turn provides data to another aggregator (not pictured) which provides data to data platform **12**). An example way to implement an aggregator is through a program written in an appropriate language, such as C or Golang.

Use of an aggregator can be beneficial in sensitive environments (e.g., involving financial or medical transactions) where various nodes are subject to regulatory or other architectural requirements (e.g., prohibiting a given node from communicating with systems outside of datacenter **104**). Use of an aggregator can also help to minimize security exposure more generally. As one example, by limiting communications with data platform **12** to data aggregator **114**, individual nodes in nodes **108** need not make external network connections (e.g., via Internet **124**), which can potentially expose them to compromise (e.g., by other external devices, such as device **118**, operated by a criminal). Similarly, data platform **12** can provide updates, configuration information, etc., to data aggregator **114**

(which in turn distributes them to nodes **108**), rather than requiring nodes **108** to allow incoming connections from data platform **12** directly.

Another benefit of an aggregator model is that network congestion can be reduced (e.g., with a single connection being made at any given time between data aggregator **114** and data platform **12**, rather than potentially many different connections being open between various of nodes **108** and data platform **12**). Similarly, network consumption can also be reduced (e.g., with the aggregator applying compression techniques/bundling data received from multiple agents).

One example way that an agent (e.g., agent **112**, installed on node **116**) can provide information to data aggregator **114** is via a REST API, formatted using data serialization protocols such as Apache Avro. One example type of information sent by agent **112** to data aggregator **114** is status information. Status information may be sent by an agent periodically (e.g., once an hour or once any other predetermined amount of time). Alternatively, status information may be sent continuously or in response to occurrence of one or more events. The status information may include, but is not limited to, a. an amount of event backlog (in bytes) that has not yet been transmitted, b. configuration information, c. any data loss period for which data was dropped, d. a cumulative count of errors encountered since the agent started, e. version information for the agent binary, and/or f. cumulative statistics on data collection (e.g., number of network packets processed, new processes seen, etc.).

A second example type of information that may be sent by agent **112** to data aggregator **114** is event data (described in more detail herein), which may include a UTC timestamp for each event. As applicable, the agent can control the amount of data that it sends to the data aggregator in each call (e.g., a maximum of 10 MB) by adjusting the amount of data sent to manage the conflicting goals of transmitting data as soon as possible, and maximizing throughput. Data can also be compressed or uncompressed by the agent (as applicable) prior to sending the data.

Each data aggregator may run within a particular customer environment. A data aggregator (e.g., data aggregator **114**) may facilitate data routing from many different agents (e.g., agents executing on nodes **108**) to data platform **12**. In various embodiments, data aggregator **114** may implement a SOCKS **5** caching proxy through which agents can connect to data platform **12**. As applicable, data aggregator **114** can encrypt (or otherwise obfuscate) sensitive information prior to transmitting it to data platform **12**, and can also distribute key material to agents which can encrypt the information (as applicable). Data aggregator **114** may include a local storage, to which agents can upload data (e.g., pcap packets). The storage may have a key-value interface. The local storage can also be omitted, and agents configured to upload data to a cloud storage or other storage area, as applicable. Data aggregator **114** can, in some embodiments, also cache locally and distribute software upgrades, patches, or configuration information (e.g., as received from data platform **12**).

Various examples associated with agent data collection and reporting will now be described.

In the following example, suppose that a user (e.g., a network administrator) at entity A (hereinafter "user A") has decided to begin using the services of data platform **12**. In some embodiments, user A may access a web frontend (e.g., web app **120**) using a computer **126** and enrolls (on behalf of entity A) an account with data platform **12**. After enrollment is complete, user A may be presented with a set of installers, pre-built and customized for the environment of

11

entity A, that user A can download from data platform 12 and deploy on nodes 108. Examples of such installers include, but are not limited to, a Windows executable file, an iOS app, a Linux package (e.g., .deb or .rpm), a binary, or a container (e.g., a Docker container). When a user (e.g., a network administrator) at entity B (hereinafter “user B”) also signs up for the services of data platform 12, user B may be similarly presented with a set of installers that are pre-built and customized for the environment of entity B.

User A deploys an appropriate installer on each of nodes 108 (e.g., with a Windows executable file deployed on a Windows-based platform or a Linux package deployed on a Linux platform, as applicable). As applicable, the agent can be deployed in a container. Agent deployment can also be performed using one or more appropriate automation tools, such as Chef, Puppet, Salt, and Ansible. Deployment can also be performed using managed/hosted container management/orchestration frameworks such as Kubernetes, Mesos, and/or Docker Swarm.

In various embodiments, the agent may be installed in the user space (i.e., is not a kernel module), and the same binary is executed on each node of the same type (e.g., all Windows-based platforms have the same Windows-based binary installed on them). An illustrative function of an agent, such as agent 112, is to collect data (e.g., associated with node 116) and report it (e.g., to data aggregator 114). Other tasks that can be performed by agents include data configuration and upgrading.

One approach to collecting data as described herein is to collect virtually all information available about a node (and, e.g., the processes running on it). Alternatively, the agent may monitor for network connections, and then begin collecting information about processes associated with the network connections, using the presence of a network packet associated with a process as a trigger for collecting additional information about the process. As an example, if a user of node 116 executes an application, such as a calculator application, which does not typically interact with the network, no information about use of that application may be collected by agent 112 and/or sent to data aggregator 114. If, however, the user of node 116 executes an ssh command (e.g., to ssh from node 116 to node 122), agent 112 may collect information about the process and provide associated information to data aggregator 114. In various embodiments, the agent may always collect/report information about certain events, such as privilege escalation, irrespective of whether the event is associated with network activity.

An approach to collecting information (e.g., by an agent) is as follows, and described in conjunction with process 200 depicted in FIG. 2A. An agent (e.g., agent 112) monitors its node (e.g., node 116) for network activity. One example way that agent 112 can monitor node 116 for network activity is by using a network packet capture tool (e.g., listening using libpcap). As packets are received (201), the agent obtains and maintains (e.g., in an in-memory cache) connection information associated with the network activity (202). Examples of such information include DNS query/response, TCP, UDP, and IP information.

The agent may also determine a process associated with the network connection (203). One example approach is for the agent to use a kernel network diagnostic API (e.g., netlink_diag) to obtain inode/process information from the kernel. Another example approach is for the agent to scan using netstat (e.g., on /proc/net/tcp, /proc/net/tcp6, /proc/net/udp, and /proc/net/udp6) to obtain sockets and relate

12

them to processes. Information such as socket state (e.g., whether a socket is connected, listening, etc.) can also be collected by the agent.

One way an agent can obtain a mapping between a given inode and a process identifier is to scan within the /proc/pid directory. For each of the processes currently running, the agent examines each of their file descriptors. If a file descriptor is a match for the inode, the agent can determine that the process associated with the file descriptor owns the inode. Once a mapping is determined between an inode and a process identifier, the mapping is cached. As additional packets are received for the connection, the cached process information is used (rather than a new search being performed).

In some cases, exhaustively scanning for an inode match across every file descriptor may not be feasible (e.g., due to CPU limitations). In various embodiments, searching through file descriptors is accordingly optimized. User filtering is one example of such an optimization. A given socket is owned by a user. Any processes associated with the socket will be owned by the same user as the socket. When matching an inode (identified as relating to a given socket) against processes, the agent can filter through the processes and only examine the file descriptors of processes sharing the same user owner as the socket. In various embodiments, processes owned by root are always searched against (e.g., even when user filtering is employed).

Another example of an optimization is to prioritize searching the file descriptors of certain processes over others. One such prioritization is to search through the subdirectories of /proc/ starting with the youngest process. One approximation of such a sort order is to search through /proc/ in reverse order (e.g., examining highest numbered processes first). Higher numbered processes are more likely to be newer (i.e., not long-standing processes), and thus more likely to be associated with new connections (i.e., ones for which inode-process mappings are not already cached). In some cases, the most recently created process may not have the highest process identifier (e.g., due to the kernel wrapping through process identifiers).

Another example prioritization is to query the kernel for an identification of the most recently created process and to search in a backward order through the directories in /proc/ (e.g., starting at the most recently created process and working backwards, then wrapping to the highest value (e.g., 32768) and continuing to work backward from there). An alternate approach is for the agent to keep track of the newest process that it has reported information on (e.g., to data aggregator 114), and begin its search of /proc/ in a forward order starting from the PID of that process.

Another example prioritization is to maintain, for each user actively using node 116, a list of the five (or any other number) most recently active processes. Those processes are more likely than other processes (less active, or passive) on node 116 to be involved with new connections, and can thus be searched first. For many processes, lower valued file descriptors tend to correspond to non-sockets (e.g., stdin, stdout, stderr). Yet another optimization is to preferentially search higher valued file descriptors (e.g., across processes) over lower valued file descriptors (that are less likely to yield matches).

In some cases, while attempting to locate a process identifier for a given inode, an agent may encounter a socket that does not correspond to the inode being matched against and is not already cached. The identity of that socket (and its corresponding inode) can be cached, once discovered, thus removing a future need to search for that pair.

13

In some cases, a connection may terminate before the agent is able to determine its associated process (e.g., due to a very short-lived connection, due to a backlog in agent processing, etc.). One approach to addressing such a situation is to asynchronously collect information about the connection using the audit kernel API, which streams information to user space. The information collected from the audit API (which can include PID/inode information) can be matched by the agent against pcap/inode information. In some embodiments, the audit API is always used, for all connections. However, due to CPU utilization considerations, use of the audit API can also be reserved for short/otherwise problematic connections (and/or omitted, as applicable).

Once the agent has determined which process is associated with the network connection (203), the agent can then collect additional information associated with the process (204). As will be described in more detail below, some of the collected information may include attributes of the process (e.g., a process parent hierarchy, and an identification of a binary associated with the process). As will also be described in more detail below, other of the collected information is derived (e.g., session summarization data and hash values).

The collected information is then transmitted (205), e.g., by an agent (e.g., agent 112) to a data aggregator (e.g., data aggregator 114), which in turn provides the information to data platform 12. In some embodiments, all information collected by an agent may be transmitted (e.g., to a data aggregator and/or to data platform 12). In other embodiments, the amount of data transmitted may be minimized (e.g., for efficiency reasons), using various techniques.

One approach to minimizing the amount of data flowing from agents (such as agents installed on nodes 108) to data platform 12 is to use a technique of implicit references with unique keys. The keys can be explicitly used by data platform 12 to extract/derive relationships, as necessary, in a data set at a later time, without impacting performance.

As previously mentioned, some data collected about a process is constant and does not change over the lifetime of the process (e.g., attributes), and some data changes (e.g., statistical information and other variable information). Constant data can be transmitted (205) once, when the agent first becomes aware of the process. And, if any changes to the constant data are detected (e.g., a process changes its parent), a refreshed version of the data can be transmitted (205) as applicable.

In some examples, an agent may collect variable data (e.g., data that may change over the lifetime of the process). In some examples, variable data can be transmitted (205) at periodic (or other) intervals. Alternatively, variable data may be transmitted in substantially real time as it is collected. In some examples, the variable data may indicate a thread count for a process, a total virtual memory used by the process, the total resident memory used by the process, the total time spent by the process executing in user space, and/or the total time spent by the process executing in kernel space. In some examples, the data may include a hash that may be used within data platform 12 to join process creation time attributes with runtime attributes to construct a full dataset.

Below are additional examples of data that an agent, such as agent 112, can collect and provide to data platform 12.

1. User Data

Core User Data: user name, UID (user ID), primary group, other groups, home directory.

14

Failed Login Data: IP address, hostname, username, count.

User Login Data: user name, hostname, IP address, start time, TTY (terminal), UID (user ID), GID (group ID), process, end time.

2. Machine Data

Dropped Packet Data: source IP address, destination IP address, destination port, protocol, count.

Machine Data: hostname, domain name, architecture, kernel, kernel release, kernel version, OS, OS version, OS description, CPU, memory, model number, number of cores, last boot time, last boot reason, tags (e.g., Cloud provider tags such as AWS, GCP, or Azure tags), default router, interface name, interface hardware address, interface IP address and mask, promiscuous mode.

3. Network Data

Network Connection Data: source IP address, destination IP address, source port, destination port, protocol, start time, end time, incoming and outgoing bytes, source process, destination process, direction of connection, histograms of packet length, inter packet delay, session lengths, etc.

Listening Ports in Server: source IP address, port number, protocol, process.

Dropped Packet Data: source IP address, destination IP address, destination port, protocol, count.

Arp Data: source hardware address, source IP address, destination hardware address, destination IP address.

DNS Data: source IP address, response code, response string, question (request), packet length, final answer (response).

4. Application Data

Package Data: exe path, package name, architecture, version, package path, checksums (MD5, SHA-1, SHA-256), size, owner, owner ID.

Application Data: command line, PID (process ID), start time, UID (user ID), EUID (effective UID), PPID (parent process ID), PGID (process group ID), SID (session ID), exe path, username, container ID.

5. Container Data

Container Image Data: image creation time, parent ID, author, container type, repo, (AWS) tags, size, virtual size, image version.

Container Data: container start time, container type, container name, container ID, network mode, privileged, PID mode, IP addresses, listening ports, volume map, process ID.

6. File Data

File path, file data hash, symbolic links, file creation data, file change data, file metadata, file mode.

As mentioned above, an agent, such as agent 112, can be deployed in a container (e.g., a Docker container), and can also be used to collect information about containers. Collection about a container can be performed by an agent irrespective of whether the agent is itself deployed in a container or not (as the agent can be deployed in a container running in a privileged mode that allows for monitoring).

Agents can discover containers (e.g., for monitoring) by listening for container create events (e.g., provided by Docker), and can also perform periodic ordered discovery scans to determine whether containers are running on a node. When a container is discovered, the agent can obtain attributes of the container, e.g., using standard Docker API calls (e.g., to obtain IP addresses associated with the container, whether there's a server running inside, what port it is listening on, associated PIDs, etc.). Information such as the parent process that started the container can also be collected, as can information about the image (which comes from the Docker repository).

15

In various embodiments, agents may use namespaces to determine whether a process is associated with a container. Namespaces are a feature of the Linux kernel that can be used to isolate resources of a collection of processes. Examples of namespaces include process ID (PID) namespaces, network namespaces, and user namespaces. Given a process, the agent can perform a fast lookup to determine whether the process is part of the namespace the container claims to be its namespace.

As mentioned, agents can be configured to report certain types of information (e.g., attribute information) once, when the agent first becomes aware of a process. In various embodiments, such static information is not reported again (or is reported once a day, every twelve hours, etc.), unless it changes (e.g., a process changes its parent, changes its owner, or a SHA-1 of the binary associated with the process changes).

In contrast to static/attribute information, certain types of data change constantly (e.g., network-related data). In various embodiments, agents are configured to report a list of current connections every minute (or other appropriate time interval). In that connection list will be connections that started in that minute interval, connections that ended in that minute interval, and connections that were ongoing throughout the minute interval (e.g., a one minute slice of a one hour connection).

In various embodiments, agents are configured to collect/compute statistical information about connections (e.g., at the one minute level of granularity and or at any other time interval). Examples of such information include, for the time interval, the number of bytes transferred, and in which direction. Another example of information collected by an agent about a connection is the length of time between packets. For connections that span multiple time intervals (e.g., a seven minute connection), statistics may be calculated for each minute of the connection. Such statistical information (for all connections) can be reported (e.g., to a data aggregator) once a minute.

In various embodiments, agents are also configured to maintain histogram data for a given network connection, and provide the histogram data (e.g., in the Apache Avro data exchange format) under the Connection event type data. Examples of such histograms include: 1. a packet length histogram (packet_len_hist), which characterizes network packet distribution; 2. a session length histogram (session_len_hist), which characterizes a network session length; 3. a session time histogram (session_time_hist), which characterizes a network session time; and 4. a session switch time histogram (session_switch_time_hist), which characterizes network session switch time (i.e., incoming→outgoing and vice versa). For example, histogram data may include one or more of the following fields: 1. count, which provides a count of the elements in the sampling; 2. sum, which provides a sum of elements in the sampling; 3. max, which provides the highest value element in the sampling; 4. std_dev, which provides the standard deviation of elements in the sampling; and 5. buckets, which provides a discrete sample bucket distribution of sampling data (if applicable).

For some protocols (e.g., HTTP), typically, a connection is opened, a string is sent, a string is received, and the connection is closed. For other protocols (e.g., NFS), both sides of the connection engage in a constant chatter. Histograms allow data platform 12 to model application behavior (e.g., using machine learning techniques), for establishing baselines, and for detecting deviations. As one example, suppose that a given HTTP server typically sends/receives 1,000 bytes (in each direction) whenever a connection is

16

made with it. If a connection generates 500 bytes of traffic, or 2,000 bytes of traffic, such connections would be considered within the typical usage pattern of the server. Suppose, however, that a connection is made that results in 10G of traffic. Such a connection is anomalous and can be flagged accordingly.

Returning to FIG. 1D, as previously mentioned, data aggregator 114 may be configured to provide information (e.g., collected from nodes 108 by agents) to data platform 12. Data aggregator 128 may be similarly configured to provide information to data platform 12. As shown in FIG. 1D, both aggregator 114 and aggregator 128 may connect to a load balancer 130, which accepts connections from aggregators (and/or as applicable, agents), as well as other devices, such as computer 126 (e.g., when it communicates with web app 120), and supports fair balancing. In various embodiments, load balancer 130 is a reverse proxy that load balances accepted connections internally to various microservices (described in more detail below), allowing for services provided by data platform 12 to scale up as more agents are added to the environment and/or as more entities subscribe to services provided by data platform 12. Example ways to implement load balancer 130 include, but are not limited to, using HaProxy, using nginx, and using elastic load balancing (ELB) services made available by Amazon.

Agent service 132 is a microservice that is responsible for accepting data collected from agents (e.g., provided by aggregator 114). In various embodiments, agent service 132 uses a standard secure protocol, such as HTTPS to communicate with aggregators (and, as applicable, agents), and receives data in an appropriate format such as Apache Avro. When agent service 132 receives an incoming connection, it can perform a variety of checks, such as to see whether the data is being provided by a current customer, and whether the data is being provided in an appropriate format. If the data is not appropriately formatted (and/or is not provided by a current customer), it may be rejected.

If the data is appropriately formatted, agent service 132 may facilitate copying the received data to a streaming data stable storage using a streaming service (e.g., Amazon Kinesis and/or any other suitable streaming service). Once the ingesting into the streaming service is complete, agent service 132 may send an acknowledgement to the data provider (e.g., data aggregator 114). If the agent does not receive such an acknowledgement, it is configured to retry sending the data to data platform 12. One way to implement agent service 132 is as a REST API server framework (e.g., Java DropWizard), configured to communicate with Kinesis (e.g., using a Kinesis library).

In various embodiments, data platform 12 uses one or more streams (e.g., Kinesis streams) for all incoming customer data (e.g., including data provided by data aggregator 114 and data aggregator 128), and the data is sharded based on the node (also referred to herein as a “machine”) that originated the data (e.g., node 116 vs. node 122), with each node having a globally unique identifier within data platform 12. Multiple instances of agent service 132 can write to multiple shards.

Kinesis is a streaming service with a limited period (e.g., 1-7 days). To persist data longer than a day, the data may be copied to long term storage 42 (e.g., S3). Data loader 136 is a microservice that is responsible for picking up data from a data stream (e.g., a Kinesis stream) and persisting it in long term storage 42. In one example embodiment, files collected by data loader 136 from the Kinesis stream are placed into one or more buckets, and segmented using a combination of a customer identifier and time slice. Given a particular time

17

segment, and a given customer identifier, the corresponding file (stored in long term storage) contains five minutes (or another appropriate time slice) of data collected at that specific customer from all of the customer's nodes. Data loader **136** can be implemented in any appropriate programming language, such as Java or C, and can be configured to use a Kinesis library to interface with Kinesis. In various embodiments, data loader **136** uses the Amazon Simple Queue Service (SQS) (e.g., to alert DB loader **140** that there is work for it to do).

DB loader **140** is a microservice that is responsible for loading data into an appropriate data store **30**, such as SnowflakeDB or Amazon Redshift, using individual per-customer databases. In particular, DB loader **140** is configured to periodically load data into a set of raw tables from files created by data loader **136** as per above. DB loader **140** manages throughput, errors, etc., to make sure that data is loaded consistently and continuously. Further, DB loader **140** can read incoming data and load into data store **30** data that is not already present in tables of data store **30** (also referred to herein as a database). DB loader **140** can be implemented in any appropriate programming language, such as Java or C, and an SQL framework such as jOOQ (e.g., to manage SQLs for insertion of data), and SQL/JDBC libraries. In some examples, DB loader **140** may use Amazon S3 and Amazon Simple Queue Service (SQS) to manage files being transferred to and from data store **30**.

Customer data included in data store **30** can be augmented with data from additional data sources, such as AWS CloudTrail and/or other types of external tracking services. To this end, data platform may include a tracking service analyzer **144**, which is another microservice. Tracking service analyzer **144** may pull data from an external tracking service (e.g., Amazon CloudTrail) for each applicable customer account, as soon as the data is available. Tracking service analyzer **144** may normalize the tracking data as applicable, so that it can be inserted into data store **30** for later querying/analysis. Tracking service analyzer **144** can be written in any appropriate programming language, such as Java or C. Tracking service analyzer **144** also makes use of SQL/JDBC libraries to interact with data store **30** to insert/query data.

As described herein, data platform **12** can model activities that occur within datacenters, such as datacenters **104** and **106**. The model may be stable over time, and differences, even subtle ones (e.g., between a current state of the datacenter and the model) can be surfaced. The ability to surface such anomalies can be particularly beneficial in datacenter environments where rogue employees and/or external attackers may operate slowly (e.g., over a period of months), hoping that the elastic nature of typical resource use (e.g., virtualized servers) will help conceal their nefarious activities.

Using techniques described herein, data platform **12** can automatically discover entities (which may implement compute assets **16**) deployed in a given datacenter. Examples of entities include workloads, applications, processes, machines, virtual machines, containers, files, IP addresses, domain names, and users. The entities may be grouped together logically (into analysis groups) based on behaviors, and temporal behavior baselines can be established. In particular, using techniques described herein, periodic graphs can be constructed (also referred to herein as poly-graphs), in which the nodes are applicable logical entities, and the edges represent behavioral relationships between the logical entities in the graph. Baselines can be created for every node and edge.

18

Communication (e.g., between applications/nodes) is one example of a behavior. A model of communications between processes is an example of a behavioral model. As another example, the launching of applications is another example of a behavior that can be modeled. The baselines may be periodically updated (e.g., hourly) for every entity. Additionally or alternatively, the baselines may be continuously updated in substantially real-time as data is collected by agents. Deviations from the expected normal behavior can then be detected and automatically reported (e.g., as anomalies or threats detected). Such deviations may be due to a desired change, a misconfiguration, or malicious activity. As applicable, data platform **12** can score the detected deviations (e.g., based on severity and threat posed). Additional examples of analysis groups include models of machine communications, models of privilege changes, and models of insider behaviors (monitoring the interactive behavior of human users as they operate within the datacenter).

Two example types of information collected by agents are network level information and process level information. As previously mentioned, agents may collect information about every connection involving their respective nodes. And, for each connection, information about both the server and the client may be collected (e.g., using the connection-to-process identification techniques described above). DNS queries and responses may also be collected. The DNS query information can be used in logical entity graphing (e.g., collapsing many different IP addresses to a single service—e.g., s3.amazonaws.com). Examples of process level information collected by agents include attributes (user ID, effective user ID, and command line). Information such as what user/application is responsible for launching a given process and the binary being executed (and its SHA-256 values) may also be provided by agents.

The dataset collected by agents across a datacenter can be very large, and many resources (e.g., virtual machines, IP addresses, etc.) are recycled very quickly. For example, an IP address and port number used at a first point in time by a first process on a first virtual machine may very rapidly be used (e.g., an hour later) by a different process/virtual machine.

A dataset (and elements within it) can be considered at both a physical level, and a logical level, as illustrated in FIG. 2B. In particular, FIG. 2B illustrates an example 5-tuple of data **210** collected by an agent, represented physically (**216**) and logically (**217**). The 5-tuple includes a source address **211**, a source port **212**, a destination address **213**, a destination port **214**, and a protocol **215**. In some cases, port numbers (e.g., **212**, **214**) may be indicative of the nature of a connection (e.g., with certain port usage standardized). However, in many cases, and in particular in datacenters, port usage is ephemeral. For example, a Docker container can listen on an ephemeral port, which is unrelated to the service it will run. When another Docker container starts (for the same service), the port may well be different. Similarly, particularly in a virtualized environment, IP addresses may be recycled frequently (and are thus also potentially ephemeral) or could be NATed, which makes identification difficult.

A physical representation of the 5-tuple is depicted in region **216**. A process **218** (executing on machine **219**) has opened a connection to machine **220**. In particular, process **218** is in communication with process **221**. Information such as the number of packets exchanged between the two machines over the respective ports can be recorded.

As previously mentioned, in a datacenter environment, portions of the 5-tuple may change-potentially frequently—

but still be associated with the same behavior. Namely, one application (e.g., Apache) may frequently be in communication with another application (e.g., Oracle), using ephemeral datacenter resources. Further, either/both of Apache and Oracle may be multi-homed. This can lead to potentially thousands of 5-tuples (or more) that all correspond to Apache communicating with Oracle within a datacenter. For example, Apache could be executed on a single machine, and could also be executed across fifty machines, which are variously spun up and down (with different IP addresses each time). An alternate representation of the 5-tuple of data **210** is depicted in region **217**, and is logical. The logical representation of the 5-tuple aggregates the 5-tuple (along with other connections between Apache and Oracle having other 5-tuples) as logically representing the same connection. By aggregating data from raw physical connection information into logical connection information, using techniques described herein, a size reduction of six orders of magnitude in the data set can be achieved.

FIG. 2C depicts a portion of a logical polygraph. Suppose a datacenter has seven instances of the application `update_engine` **225**, executing as seven different processes on seven different machines, having seven different IP addresses, and using seven different ports. The instances of `update_engine` variously communicate with `update.core-os.net` **226**, which may have a single IP address or many IP addresses itself, over the one hour time period represented in the polygraph. In the example shown in FIG. 2C, `update_engine` is a client, connecting to the server `update.core-os.net`, as indicated by arrow **228**.

Behaviors of the seven processes are clustered together, into a single summary. As indicated in region **227**, statistical information about the connections is also maintained (e.g., number of connections, histogram information, etc.). A polygraph such as is depicted in FIG. 2C can be used to establish a baseline of behavior (e.g., at the one-hour level), allowing for the future detection of deviations from that baseline. As one example, suppose that statistically an `update_engine` instance transmits data at 11 bytes per second. If an instance were instead to transmit data at 1000 bytes per second, such behavior would represent a deviation from the baseline and could be flagged accordingly. Similarly, changes that are within the baseline (e.g., an eighth instance of `update_engine` appears, but otherwise behaves as the other instances; or one of the seven instances disappears) are not flagged as anomalous. Further, datacenter events, such as failover, autobalancing, and A-B refresh are unlikely to trigger false alarms in a polygraph, as at the logical level, the behaviors remain the same.

In various embodiments, polygraph data is maintained for every application in a datacenter, and such polygraph data can be combined to make a single datacenter view across all such applications. FIG. 2D illustrates a portion of a polygraph for a service that evidences more complex behaviors than are depicted in FIG. 2C. In particular, FIG. 2D illustrates the behaviors of S3 as a service (as used by a particular customer datacenter). Clients within the datacenter variously connect to the S3 service using one of five fully qualified domains (listed in region **230**). Contact with any of the domains is aggregated as contact with S3 (as indicated in region **231**). Depicted in region **232** are various containers which (as clients) connect with S3. Other containers (which do not connect with S3) are not included. As with the polygraph portion depicted in FIG. 2C, statistical information about the connections is known and summarized, such as the number of bytes transferred, histogram information, etc.

FIG. 2E illustrates a communication polygraph for a datacenter. In particular, the polygraph indicates a one hour summary of approximately 500 virtual machines, which collectively run one million processes, and make 100 million connections in that hour. As illustrated in FIG. 2E, a polygraph represents a drastic reduction in size (e.g., from tracking information on 100 million connections in an hour, to a few hundred nodes and a few hundred edges). Further, as a datacenter scales up (e.g., from using 10 virtual machines to 100 virtual machines as the datacenter uses more workers to support existing applications), the polygraph for the datacenter will tend to stay the same size (with the 100 virtual machines clustering into the same nodes that the 10 virtual machines previously clustered into). As new applications are added into the datacenter, the polygraph may automatically scale to include behaviors involving those applications.

In the particular polygraph shown in FIG. 2E, nodes generally correspond to workers, and edges correspond to communications the workers engage in (with connection activity being the behavior modeled in polygraph **235**). Another example polygraph could model other behavior, such as application launching. The communications graphed in FIG. 2E include traffic entering the datacenter, traffic exiting the datacenter, and traffic that stays wholly within the datacenter (e.g., traffic between workers). One example of a node included in polygraph **235** is the `sshd` application, depicted as node **236**. As indicated in FIG. 2E, **421** instances of `sshd` were executing during the one hour time period of data represented in polygraph **235**. As indicated in region **237**, nodes within the datacenter communicated with a total of 1349 IP addresses outside of the datacenter (and not otherwise accounted for, e.g., as belonging to a service such as Amazon AWS **238** or Slack **239**).

In the following examples, suppose that user B, an administrator of datacenter **106**, is interacting with data platform **12** to view visualizations of polygraphs in a web browser (e.g., as served to user B via web app **120**). One type of polygraph user B can view is an application-communication polygraph, which indicates, for a given one hour window (or any other suitable time interval), which applications communicated with which other applications. Another type of polygraph user B can view is an application launch polygraph. User B can also view graphs related to user behavior, such as an insider behavior graph which tracks user connections (e.g., to internal and external applications, including chains of such behavior), a privilege change graph which tracks how privileges change between processes, and a user login graph, which tracks which (logical) machines a user logs into.

FIG. 2F illustrates an example of an application-communication polygraph for a datacenter (e.g., datacenter **106**) for the one hour period of 9 am-10 am on June 5. The time slice currently being viewed is indicated in region **240**. If user B clicks his mouse in region **241**, user B will be shown a representation of the application-communication polygraph as generated for the following hour (10 am-11 am on June 5).

FIG. 2G depicts what is shown in user B's browser after he has clicked on region **241**, and has further clicked on region **242**. The selection in region **242** turns on and off the ability to compare two time intervals to one another. User B can select from a variety of options when comparing the 9 am-10 am and 10 am-11 am time intervals. By clicking region **248**, user B will be shown the union of both graphs (i.e., any connections that were present in either time interval). By clicking region **249**, user B will be shown the

21

intersection of both graphs (i.e., only those connections that were present in both time intervals).

As shown in FIG. 2G, user B has elected to click on region 250, which depicts connections that are only present in the 9 am-10 am polygraph in a first color 251, and depicts connections that are only present in the 10 am-11 am polygraph in a second color 252. Connections present in both polygraphs are omitted from display. As one example, in the 9 am-10 am polygraph (corresponding to connections made during the 9 am-10 am time period at datacenter 106), a connection was made by a server to sshd (253) and also to systemd (254). Both of those connections ended prior to 10 am and are thus depicted in the first color. As another example, in the 10 am-11 am polygraph (corresponding to connections made during the 10 am-11 am time period at datacenter 106), a connection was made from a known bad external IP to nginx (255). The connection was not present during the 9 am-10 am time slice and thus is depicted in the second color. As yet another example, two different connections were made to a Slack service between 9 am and 11 am. However, the first was made by a first client during the 9 am-10 am time slice (256) and the second was made by a different client during the 10 am-11 am slice (257), and so the two connections are depicted respectively in the first and second colors and blue.

Returning to the polygraph depicted in FIG. 2F, suppose user B enters “etcd” into the search box located in region 244. User B will then be presented with the interface illustrated in FIG. 2H. As shown in FIG. 2H, three applications containing the term “etcd” were engaged in communications during the 9 am-10 am window. One application is etcdctl, a command line client for etcd. As shown in FIG. 2H, a total of three different etcdctl processes were executed during the 9 am-10 am window, and were clustered together (260). FIG. 2H also depicts two different clusters that are both named etcd2. The first cluster includes (for the 9 am-10 am window) five members (261) and the second cluster includes (for the same window) eight members (262). The reason for these two distinct clusters is that the two groups of applications behave differently (e.g., they exhibit two distinct sets of communication patterns). Specifically, the instances of etcd2 in cluster 261 only communicate with locksmithctl (263) and other etcd2 instances (in both clusters 261 and 262). The instances of etcd2 in cluster 262 communicate with additional entities, such as etcdctl and Docker containers. As desired, user B can click on one of the clusters (e.g., cluster 261) and be presented with summary information about the applications included in the cluster, as is shown in FIG. 2I (e.g., in region 265). User B can also double click on a given cluster (e.g., cluster 261) to see details on each of the individual members of the cluster broken out.

Suppose user B now clicks on region 245 of the interface shown in FIG. 2F. User B will then be shown an application launch polygraph. Launching an application is another example of a behavior. The launch polygraph models how applications are launched by other applications. FIG. 2J illustrates an example of a portion of a launch polygraph. In particular, user B has typed “find” into region 266, to see how the “find” application is being launched. As shown in FIG. 2J, in the launch polygraph for the 10 am-11 am time period, find applications (267) are always launched by bash (268), which is in turn always launched by systemd (269). If find is launched by a different application, this would be anomalous behavior.

FIG. 2K illustrates another example of a portion of an application launch polygraph. In FIG. 2K, user B has

22

searched (270) for “python ma” to see how “python marathon_lb” (271) is launched. As shown in FIG. 2K, in each case (during the one hour time slice of 10 am-11 am), python marathon_lb is launched as a result of a chain of the same seven applications each time. If python marathon_lb is ever launched in a different manner, this indicates anomalous behavior. The behavior could be indicative of malicious activities, but could also be due to other reasons, such as a misconfiguration, a performance-related issue, and/or a failure, etc.

Suppose user B now clicks on region 246 of the interface shown in FIG. 2F. User B will then be shown an insider behavior graph. The insider behavior graph tracks information about behaviors such as processes started by a user interactively using protocols such as ssh or telnet, and any processes started by those processes. As one example, suppose an administrator logs into a first virtual machine in datacenter 106 (e.g., using sshd via an external connection he makes from a hotel), using a first set of credentials (e.g., first.last@example.com and an appropriate password). From the first virtual machine, the administrator connects to a second virtual machine (e.g., using the same credentials), then uses the sudo command to change identities to those of another user, and then launches a program. Graphs built by data platform 12 can be used to associate the administrator with each of his actions, including launching the program using the identity of another user.

FIG. 2L illustrates an example of a portion of an insider behavior graph. In particular, in FIG. 2L, user B is viewing a graph that corresponds to the time slice of 3 pm-4 pm on June 1. FIG. 2L illustrates the internal/external applications that users connected to during the one hour time slice. If a user typically communicates with particular applications, that information will become part of a baseline. If the user deviates from his baseline behavior (e.g., using new applications, or changing privilege in anomalous ways), such anomalies can be surfaced.

FIG. 2M illustrates an example of a portion of a privilege change graph, which identifies how privileges are changed between processes. Typically, when a user launches a process (e.g., “ls”), the process inherits the same privileges that the user has. And, while a process can have fewer privileges than the user (i.e., go down in privilege), it is rare (and generally undesirable) for a user to escalate in privilege. Information included in the privilege change graph can be determined by examining the parent of each running process, and determining whether there is a match in privilege between the parent and the child. If the privileges are different, a privilege change has occurred (whether a change up or a change down). The application ntpd is one rare example of a scenario in which a process escalates (272) to root, and then returns back (273). The sudo command is another example (e.g., used by an administrator to temporarily have a higher privilege). As with the other examples, ntpd’s privilege change actions, and the legitimate actions of various administrators (e.g., using sudo) will be incorporated into a baseline model by data platform 12. When deviations occur, such as where a new application that is not ntpd escalates privilege, or where an individual that has not previously/does not routinely use sudo does so, such behaviors can be identified as anomalous.

FIG. 2N illustrates an example of a portion of a user login graph, which identifies which users log into which logical nodes. Physical nodes (whether bare metal or virtualized) are clustered into a logical machine cluster, for example, using yet another graph, a machine-server graph, an example of which is shown in FIG. 2O. For each machine, a deter-

mination is made as to what type of machine it is, based on what kind(s) of workflows it runs. As one example, some machines run as master nodes (having a typical set of workflows they run, as master nodes) and can thus be clustered as master nodes. Worker nodes are different from master nodes, for example, because they run Docker containers, and frequently change as containers move around. Worker nodes can similarly be clustered.

As previously mentioned, the polygraph depicted in FIG. 2E corresponds to activities in a datacenter in which, in a given hour, approximately 500 virtual machines collectively run one million processes, and make 100 million connections in that hour. The polygraph represents a drastic reduction in size (e.g., from tracking information on 100 million connections in an hour, to a few hundred nodes and a few hundred edges). Using techniques described herein, such a polygraph can be constructed (e.g., using commercially available computing infrastructure) in less than an hour (e.g., within a few minutes). Thus, ongoing hourly snapshots of a datacenter can be created within a two hour moving window (i.e., collecting data for the time period 8 am-9 am, while also generating a snapshot for the time previous time period 7 am-8 am). The following describes various example infrastructure that can be used in polygraph construction, and also describes various techniques that can be used to construct polygraphs.

Returning to FIG. 1D, embodiments of data platform 12 may be built using any suitable infrastructure as a service (IaaS) (e.g., AWS). For example, data platform 12 can use Simple Storage Service (S3) for data storage, Key Management Service (KMS) for managing secrets, Simple Queue Service (SQS) for managing messaging between applications, Simple Email Service (SES) for sending emails, and Route 53 for managing DNS. Other infrastructure tools can also be used. Examples include: orchestration tools (e.g., Kubernetes or Mesos/Marathon), service discovery tools (e.g., Mesos-DNS), service load balancing tools (e.g., marathon-LB), container tools (e.g., Docker or rkt), log/metric tools (e.g., collectd, fluentd, kibana, etc.), big data processing systems (e.g., Spark, Hadoop, AWS Redshift, Snowflake etc.), and distributed key value stores (e.g., Apache Zookeeper or etcd2).

As previously mentioned, in various embodiments, data platform 12 may make use of a collection of microservices. Each microservice can have multiple instances, and may be configured to recover from failure, scale, and distribute work amongst various such instances, as applicable. For example, microservices are auto-balancing for new instances, and can distribute workload if new instances are started or existing instances are terminated. In various embodiments, microservices may be deployed as self-contained Docker containers. A Mesos-Marathon or Spark framework can be used to deploy the microservices (e.g., with Marathon monitoring and restarting failed instances of microservices as needed). The service etcd2 can be used by microservice instances to discover how many peer instances are running, and used for calculating a hash-based scheme for workload distribution. Microservices may be configured to publish various health/status metrics to either an SQS queue, or etcd2, as applicable. In some examples, Amazon DynamoDB can be used for state management.

Additional information on various microservices used in embodiments of data platform 12 is provided below.

Graph generator 146 is a microservice that may be responsible for generating raw behavior graphs on a per customer basis periodically (e.g., once an hour). In particular, graph generator 146 may generate graphs of entities (as

the nodes in the graph) and activities between entities (as the edges). In various embodiments, graph generator 146 also performs other functions, such as aggregation, enrichment (e.g., geolocation and threat), reverse DNS resolution, TF-IDF based command line analysis for command type extraction, parent process tracking, etc.

Graph generator 146 may perform joins on data collected by the agents, so that both sides of a behavior are linked. For example, suppose a first process on a first virtual machine (e.g., having a first IP address) communicates with a second process on a second virtual machine (e.g., having a second IP address). Respective agents on the first and second virtual machines may each report information on their view of the communication (e.g., the PID of their respective processes, the amount of data exchanged and in which direction, etc.). When graph generator performs a join on the data provided by both agents, the graph will include a node for each of the processes, and an edge indicating communication between them (as well as other information, such as the directionality of the communication—i.e., which process acted as the server and which as the client in the communication).

In some cases, connections are process to process (e.g., from a process on one virtual machine within the cloud environment associated with entity A to another process on a virtual machine within the cloud environment associated with entity A). In other cases, a process may be in communication with a node (e.g., outside of entity A) which does not have an agent deployed upon it. As one example, a node within entity A might be in communication with node 172, outside of entity A. In such a scenario, communications with node 172 are modeled (e.g., by graph generator 146) using the IP address of node 172. Similarly, where a node within entity A does not have an agent deployed upon it, the IP address of the node can be used by graph generator in modeling.

Graphs created by graph generator 146 may be written to data store 30 and cached for further processing. A graph may be a summary of all activity that happened in a particular time interval. As each graph corresponds to a distinct period of time, different rows can be aggregated to find summary information over a larger timestamp. In some examples, picking two different graphs from two different timestamps can be used to compare different periods. If necessary, graph generator 146 can parallelize its workload (e.g., where its backlog cannot otherwise be handled within a particular time period, such as an hour, or if is required to process a graph spanning a long time period).

Graph generator 146 can be implemented in any appropriate programming language, such as Java or C, and machine learning libraries, such as Spark's MLlib. Example ways that graph generator computations can be implemented include using SQL or Map-R, using Spark or Hadoop.

SSH tracker 148 is a microservice that may be responsible for following ssh connections and process parent hierarchies to determine trails of user ssh activity. Identified ssh trails are placed by the SSH tracker 148 into data store 30 and cached for further processing.

SSH tracker 148 can be implemented in any appropriate programming language, such as Java or C, and machine learning libraries, such as Spark's MLlib. Example ways that SSH tracker computations can be implemented include using SQL or Map-R, using Spark or Hadoop.

Threat aggregator 150 is a microservice that may be responsible for obtaining third party threat information from various applicable sources, and making it available to other micro-services. Examples of such information include reverse DNS information, GeoIP information, lists of known

25

bad domains/IP addresses, lists of known bad files, etc. As applicable, the threat information is normalized before insertion into data store 30. Threat aggregator 150 can be implemented in any appropriate programming language, such as Java or C, using SQL/JDBC libraries to interact with data store 30 (e.g., for insertions and queries).

Scheduler 152 is a microservice that may act as a scheduler and that may run arbitrary jobs organized as a directed graph. In some examples, scheduler 152 ensures that all jobs for all customers are able to run during a given time interval (e.g., every hour). Scheduler 152 may handle errors and retrying for failed jobs, track dependencies, manage appropriate resource levels, and/or scale jobs as needed. Scheduler 152 can be implemented in any appropriate programming language, such as Java or C. A variety of components can also be used, such as open source scheduler frameworks (e.g., Airflow), or AWS services (e.g., the AWS Data pipeline) which can be used for managing schedules.

Graph Behavior Modeler (GBM) 154 is a microservice that may compute polygraphs. In particular, GBM 154 can be used to find clusters of nodes in a graph that should be considered similar based on some set of their properties and relationships to other nodes. As described herein, the clusters and their relationships can be used to provide visibility into a datacenter environment without requiring user specified labels. GBM 154 may track such clusters over time persistently, allowing for changes to be detected and alerts to be generated.

GBM 154 may take as input a raw graph (e.g., as generated by graph generator 146). Nodes are actors of a behavior, and edges are the behavior relationship itself. For example, in the case of communication, example actors include processes, which communicate with other processes. The GBM 154 clusters the raw graph based on behaviors of actors and produces a summary (the polygraph). The polygraph summarizes behavior at a datacenter level. The GBM 154 also produces “observations” that represent changes detected in the datacenter. Such observations may be based on differences in cumulative behavior (e.g., the baseline) of the datacenter with its current behavior. The GBM 154 can be implemented in any appropriate programming language, such as Java, C, or Golang, using appropriate libraries (as applicable) to handle distributed graph computations (handling large amounts of data analysis in a short amount of time). Apache Spark is another example tool that can be used to compute polygraphs. The GBM 154 can also take feedback from users and adjust the model according to that feedback. For example, if a given user is interested in relearning behavior for a particular entity, the GBM 154 can be instructed to “forget” the implicated part of the polygraph.

GBM runner 156 is a microservice that may be responsible for interfacing with GBM 154 and providing GBM 154 with raw graphs (e.g., using a query language, such as SQL, to push any computations it can to data store 30). GBM runner 156 may also insert polygraph output from GBM 154 to data store 30. GBM runner 156 can be implemented in any appropriate programming language, such as Java or C, using SQL/JDBC libraries to interact with data store 30 to insert and query data.

Alert generator 158 is a microservice that may be responsible for generating alerts. Alert generator 158 may examine observations (e.g., produced by GBM 154) in aggregate, deduplicate them, and score them. Alerts may be generated for observations with a score exceeding a threshold. Alert generator 158 may also compute (or retrieve, as applicable) data that a customer (e.g., user A or user B) might need when

26

reviewing the alert. Examples of events that can be detected by data platform 12 (and alerted on by alert generator 158) include, but are not limited to the following:

new user: This event may be created the first time a user (e.g., of node 116) is first observed by an agent within a datacenter.

user launched new binary: This event may be generated when an interactive user launches an application for the first time.

new privilege escalation: This event may be generated when user privileges are escalated and a new application is run.

new application or container: This event may be generated when an application or container is seen for the first time.

new external connection: This event may be generated when a connection to an external IP/domain is made from a new application.

new external host or IP: This event may be generated when a new external host or IP is involved in a connection with a datacenter.

new internal connection: This event may be generated when a connection between internal-only applications is seen for the first time.

new external client: This event may be generated when a new external connection is seen for an application which typically does not have external connections.

new parent: This event may be generated when an application is launched by a different parent.

connection to known bad IP/domain: Data platform 12 maintains (or can otherwise access) one or more reputation feeds. If an environment makes a connection to a known bad IP or domain, an event will be generated.

login from a known bad IP/domain: An event may be generated when a successful connection to a datacenter from a known bad IP is observed by data platform 12.

Alert generator 158 can be implemented in any appropriate programming language, such as Java or C, using SQL/JDBC libraries to interact with data store 30 to insert and query data. In various embodiments, alert generator 158 also uses one or more machine learning libraries, such as Spark’s MLLib (e.g., to compute scoring of various observations). Alert generator 158 can also take feedback from users about which kinds of events are of interest and which to suppress.

QsJobServer 160 is a microservice that may look at all the data produced by data platform 12 for an hour, and compile a materialized view (MV) out of the data to make queries faster. The MV helps make sure that the queries customers most frequently run, and data that they search for, can be easily queried and answered. QsJobServer 160 may also precompute and cache a variety of different metrics so that they can quickly be provided as answers at query time. QsJobServer 160 can be implemented using any appropriate programming language, such as Java or C, using SQL/JDBC libraries. In some examples, QsJobServer 160 is able to compute an MV efficiently at scale, where there could be a large number of joins. An SQL engine, such as Oracle, can be used to efficiently execute the SQL, as applicable.

Alert notifier 162 is a microservice that may take alerts produced by alert generator 158 and send them to customers’ integrated Security Information and Event Management (SIEM) products (e.g., Splunk, Slack, etc.). Alert notifier 162 can be implemented using any appropriate programming language, such as Java or C. Alert notifier 162 can be configured to use an email service (e.g., AWS SES or pagerduty) to send emails. Alert notifier 162 may also

provide templating support (e.g., Velocity or Moustache) to manage templates and structured notifications to SIEM products.

Reporting module **164** is a microservice that may be responsible for creating reports out of customer data (e.g., daily summaries of events, etc.) and providing those reports to customers (e.g., via email). Reporting module **164** can be implemented using any appropriate programming language, such as Java or C. Reporting module **164** can be configured to use an email service (e.g., AWS SES or pagerduty) to send emails. Reporting module **164** may also provide templating support (e.g., Velocity or Moustache) to manage templates (e.g., for constructing HTML-based email).

Web app **120** is a microservice that provides a user interface to data collected and processed on data platform **12**. Web app **120** may provide login, authentication, query, data visualization, etc. features. Web app **120** may, in some embodiments, include both client and server elements. Example ways the server elements can be implemented are using Java

DropWizard or Node.js to serve business logic, and a combination of JSON/HTTP to manage the service. Example ways the client elements can be implemented are using frameworks such as React, Angular, or Backbone. JSON, jQuery, and JavaScript libraries (e.g., underscore) can also be used.

Query service **166** is a microservice that may manage all database access for web app **120**. Query service **166** abstracts out data obtained from data store **30** and provides a JSON-based REST API service to web app **120**. Query service **166** may generate SQL queries for the REST APIs that it receives at run time. Query service **166** can be implemented using any appropriate programming language, such as Java or C and SQL/JDBC libraries, or an SQL framework such as jOOQ. Query service **166** can internally make use of a variety of types of databases, including a relational database engine **168** (e.g., AWS Aurora) and/or data store **30** to manage data for clients. Examples of tables that query service **166** manages are OLTP tables and data warehousing tables.

Cache **170** may be implemented by Redis and/or any other service that provides a key-value store. Data platform **12** can use cache **170** to keep information for frontend services about users. Examples of such information include valid tokens for a customer, valid cookies of customers, the last time a customer tried to login, etc.

FIG. 3A illustrates an example of a process for detecting anomalies in a network environment. In various embodiments, process **300** is performed by data platform **12**. The process begins at **301** when data associated with activities occurring in a network environment (such as entity A's datacenter) is received. One example of such data that can be received at **301** is agent-collected data described above (e.g., in conjunction with process **200**).

At **302**, a logical graph model is generated, using at least a portion of the monitored activities. A variety of approaches can be used to generate such logical graph models, and a variety of logical graphs can be generated (whether using the same, or different approaches). The following is one example of how data received at **301** can be used to generate and maintain a model.

During bootstrap, data platform **12** creates an aggregate graph of physical connections (also referred to herein as an aggregated physical graph) by matching connections that occurred in the first hour into communication pairs. Clustering is then performed on the communication pairs. Examples of such clustering, described in more detail below,

include performing Matching Neighbor clustering and similarity (e.g., SimRank) clustering. Additional processing can also be performed (and is described in more detail below), such as by splitting clusters based on application type, and annotating nodes with DNS query information. The resulting graph (also referred to herein as a base graph or common graph) can be used to generate a variety of models, where a subset of node and edge types (described in more detail below) and their properties are considered in a given model. One example of a model is a UID to UID model (also referred to herein as a Uid2Uid model) which clusters together processes that share a username and show similar privilege change behavior. Another example of a model is a CType model, which clusters together processes that share command line similarity. Yet another example of a model is a PType model, which clusters together processes that share behaviors over time.

Each hour (or any other predetermined time interval) after bootstrap, a new snapshot is taken (i.e., data collected about a datacenter in the last hour is processed) and information from the new snapshot is merged with existing data to create and (as additional data is collected/processed) maintain a cumulative graph. The cumulative graph (also referred to herein as a cumulative PType graph and a polygraph) is a running model of how processes behave over time. Nodes in the cumulative graph are PType nodes, and provide information such as a list of all active processes and PIDs in the last hour, the number of historic total processes, the average number of active processes per hour, the application type of the process (e.g., the CType of the PType), and historic CType information/frequency. Edges in the cumulative graph can represent connectivity and provide information such as connectivity frequency. The edges can be weighted (e.g., based on number of connections, number of bytes exchanged, etc.). Edges in the cumulative graph (and snapshots) can also represent transitions.

One approach to merging a snapshot of the activity of the last hour into a cumulative graph is as follows. An aggregate graph of physical connections is made for the connections included in the snapshot (as was previously done for the original snapshot used during bootstrap). And, clustering/splitting is similarly performed on the snapshot's aggregate graph. Next, PType clusters in the snapshot's graph are compared against PType clusters in the cumulative graph to identify commonality.

One approach to determining commonality is, for any two nodes that are members of a given CmdType (described in more detail below), comparing internal neighbors and calculating a set membership Jaccard distance. The pairs of nodes are then ordered by decreasing similarity (i.e., with the most similar sets first). For nodes with a threshold amount of commonality (e.g., at least 66% members in common), any new nodes (i.e., appearing in the snapshot's graph but not the cumulative graph) are assigned the same PType identifier as is assigned to the corresponding node in the cumulative graph. For each node that is not classified (i.e., has not been assigned a PType identifier), a network signature is generated (i.e., indicative of the kinds of network connections the node makes, who the node communicates with, etc.). The following processing is then performed until convergence. If a match of the network signature is found in the cumulative graph, the unclassified node is assigned the PType identifier of the corresponding node in the cumulative graph. Any nodes which remain unclassified after convergence are new PTypes and are assigned new identifiers and added to the cumulative graph as new. As applicable, the detection of a new PType can be

29

used to generate an alert. If the new PType has a new CmdType, a severity of the alert can be increased. If any surviving nodes (i.e., present in both the cumulative graph and the snapshot graph) change PTypes, such change is noted as a transition, and an alert can be generated. Further, if a surviving node changes PType and also changes Cmd-Type, a severity of the alert can be increased.

Changes to the cumulative graph (e.g., a new PType or a new edge between two PTypes) can be used (e.g., at **303**) to detect anomalies (described in more detail below). Two example kinds of anomalies that can be detected by data platform **12** include security anomalies (e.g., a user or process behaving in an unexpected manner) and devops/root cause anomalies (e.g., network congestion, application failure, etc.). Detected anomalies can be recorded and surfaced (e.g., to administrators, auditors, etc.), such as through alerts which are generated at **304** based on anomaly detection.

Additional detail regarding processing performed, by various components depicted in FIG. **1D** (whether performed individually or in combination), in conjunction with model/polygraph construction (e.g., as performed at **302**) are provided below.

As explained above, an aggregated physical graph can be generated on a per customer basis periodically (e.g., once an hour) from raw physical graph information, by matching connections (e.g., between two processes on two virtual machines). In various embodiments, a deterministic fixed approach is used to cluster nodes in the aggregated physical graph (e.g., representing processes and their communications). As one example, Matching Neighbors Clustering (MNC) can be performed on the aggregated physical graph to determine which entities exhibit identical behavior and cluster such entities together.

FIG. **3B** depicts a set of example processes (p1, p2, p3, and p4) communicating with other processes (p10 and p11). FIG. **3B** is a graphical representation of a small portion of an aggregated physical graph showing (for a given time period, such as an hour) which processes in a datacenter communicate with which other processes. Using MNC, processes p1, p2, and p3 will be clustered together (**305**), as they exhibit identical behavior (they communicate with p10 and only p10). Process p4, which communicates with both p10 and p11, will be clustered separately.

In MNC, only those processes exhibiting identical (communication) behavior will be clustered. In various embodiments, an alternate clustering approach can also/instead be used, which uses a similarity measure (e.g., constrained by a threshold value, such as a 60% similarity) to cluster items. In some embodiments, the output of MNC is used as input to SimRank, in other embodiments, MNC is omitted.

FIG. **3C** depicts a set of example processes (p4, p5, p6) communicating with other processes (p7, p8, p9). As illustrated, most of nodes p4, p5, and p6 communicate with most of nodes p7, p8, and p9 (as indicated in FIG. **3C** with solid connection lines). As one example, process p4 communicates with process p7 (**310**), process p8 (**311**), and process p9 (**312**). An exception is process p6, which communicates with processes p7 and p8, but does not communicate with process p9 (as indicated by dashed line **313**). If MNC were applied to the nodes depicted in FIG. **3C**, nodes p4 and p5 would be clustered (and node p6 would not be included in their cluster).

One approach to similarity clustering is to use SimRank. In an embodiment of the SimRank approach, for a given node v in a directed graph, $I(v)$ and $O(v)$ denote the respective set of in-neighbors and out-neighbors of v . Individual in-neighbors are denoted as $I_i(v)$, for $1 \leq i \leq |I(v)|$, and

30

individual out-neighbors are denoted as $O_i(v)$, for $1 \leq i \leq |O(v)|$. The similarity between two objects a and b can be denoted by $s(a,b) \in [1,0]$. A recursive equation (hereinafter “the SimRank equation”) can be written for $s(a,b)$, where, if $a=b$, then $s(a,b)$ is defined as 1, otherwise,

$$s(a, b) = \frac{C}{|I(a)||I(b)|} \sum_{i=1}^{|I(a)|} \sum_{j=1}^{|I(b)|} s(I_i(a), I_j(b))$$

where C is a constant between 0 and 1. One example value for the decay factor C is 0.8 (and a fixed number of iterations such as five). Another example value for the decay factor C is 0.6 (and/or a different number of iterations). In the event that a or b has no in-neighbors, similarity is set to $s(a,b)=0$, so the summation is defined to be 0 when $I(a)=\emptyset$ or $I(b)=\emptyset$.

The SimRank equations for a graph G can be solved by iteration to a fixed point. Suppose n is the number of nodes in G . For each iteration k , n^2 entries $s_k(*,*)$ are kept, where $s_k(a,b)$ gives the score between a and b on iteration k . Successive computations of $s_{k+1}(*,*)$ are made based on $s_k(*,*)$. Starting with $s_0(*,*)$, where each $s_0(a,b)$ is a lower bound on the actual SimRank score

$$s(a, b) : s_0(a, b) = \begin{cases} 1, & \text{if } a = b, \\ 0, & \text{if } a \neq b. \end{cases}$$

The SimRank equation can be used to compute $s_{k+1}(a,b)$ from $s_k(*,*)$ with

$$s_{k+1}(a, b) = \frac{C}{|I(a)||I(b)|} \sum_{i=1}^{|I(a)|} \sum_{j=1}^{|I(b)|} s_k(I_i(a), I_j(b))$$

for $a \neq b$, and $s_{k+1}(a,b)=1$ for $a=b$. On each iteration $k+1$, the similarity of (a,b) is updated using the similarity scores of the neighbors of (a,b) from the previous iteration k according to the SimRank equation. The values $s_k(*,*)$ are nondecreasing as k increases.

Returning to FIG. **3C**, while MNC would cluster nodes p4 and p5 together (and not include node p6 in their cluster), application of SimRank would cluster nodes p4-p6 into one cluster (**314**) and also cluster nodes p7-p9 into another cluster (**315**).

FIG. **3D** depicts a set of processes, and in particular server processes s1 and s2, and client processes c1, c2, c3, c4, c5, and c6. Suppose only nodes s1, s2, c1, and c2 are present in the graph depicted in FIG. **3D** (and the other nodes depicted are omitted from consideration). Using MNC, nodes s1 and s2 would be clustered together, as would nodes c1 and c2. Performing SimRank clustering as described above would also result in those two clusters (s1 and s2, and c1 and c2). As previously mentioned, in MNC, identical behavior is required. Thus, if node c3 were now also present in the graph, MNC would not include c3 in a cluster with c2 and c1 because node c3 only communicates with node s2 and not node s1. In contrast, a SimRank clustering of a graph that includes nodes s1, s2, c1, c2, and c3 would result (based, e.g., on an applicable selected decay value and number of iterations) in a first cluster comprising nodes s1 and s2, and a second cluster of c1, c2, and c3. As an increasing number of nodes which communicate with server process s2, and do not also communicate with server process s1, are included in

31

the graph (e.g., as c4, c5, and c6 are added), under SimRank, nodes s1 and s2 will become decreasingly similar (i.e., their intersection is reduced).

In various embodiments, SimRank is modified (from what is described above) to accommodate differences between the asymmetry of client and server connections. As one example, SimRank can be modified to use different thresholds for client communications (e.g., an 80% match among nodes c1-c6) and for server communications (e.g., a 60% match among nodes s1 and s2). Such modification can also help achieve convergence in situations such as where a server process dies on one node and restarts on another node.

The application of MNC/SimRank to an aggregated physical graph results in a smaller graph, in which processes which are determined to be sufficiently similar are clustered together. Typically, clusters generated as output of MNC will be underinclusive. For example, for the nodes depicted in FIG. 3C, process p6 will not be included in a cluster with processes p4 and p5, despite substantial similarity in their communication behaviors. The application of SimRank (e.g., to the output of MNC) helps mitigate the underinclusiveness of MNC, but can result in overly inclusive clusters. As one example, suppose (returning to the nodes depicted in FIG. 3B) that as a result of applying SimRank to the depicted nodes, nodes p1-p4 are all included in a single cluster. Both MNC and SimRank operate agnostically of which application a given process belongs to. Suppose processes p1-p3 each correspond to a first application (e.g., an update engine), and process p4 corresponds to a second application (e.g., sshd). Further suppose process p10 corresponds to contact with AWS. Clustering all four of the processes together (e.g., as a result of SimRank) could be problematic, particularly in a security context (e.g., where granular information useful in detecting threats would be lost).

As previously mentioned, data platform 12 may maintain a mapping between processes and the applications to which they belong. In various embodiments, the output of SimRank (e.g., SimRank clusters) is split based on the applications to which cluster members belong (such a split is also referred to herein as a “CmdType split”). If all cluster members share a common application, the cluster remains. If different cluster members originate from different applications, the cluster members are split along application-type (CmdType) lines. Using the nodes depicted in FIG. 3D as an example, suppose that nodes c1, c2, c3, and c5 all share “update engine” as the type of application to which they belong (sharing a CmdType). Suppose that node c4 belongs to “ssh,” and suppose that node c6 belongs to “bash.” As a result of SimRank, all six nodes (c1-c6) might be clustered into a single cluster. After a CmdType split is performed on the cluster, however, the single cluster will be broken into three clusters (c1, c2, c3, c5; c4; and c6). Specifically, the resulting clusters comprise processes associated with the same type of application, which exhibit similar behaviors (e.g., communication behaviors). Each of the three clusters resulting from the CmdType split represents, respectively, a node (also referred to herein as a PType) of a particular CmdType. Each PType is given a persistent identifier and stored persistently as a cumulative graph.

A variety of approaches can be used to determine a CmdType for a given process. As one example, for some applications (e.g., sshd), a one-to-one mapping exists between the CmdType and the application/binary name. Thus, processes corresponding to the execution of sshd will be classified using a CmdType of sshd. In various embodiments, a list of common application/binary names (e.g., sshd, apache, etc.) is maintained by data platform 12 and

32

manually curated as applicable. Other types of applications (e.g., Java, Python, and Ruby) are multi-homed, meaning that several very different applications may all execute using the binary name, “java.” For these types of applications, information such as command line/execution path information can be used in determining a CmdType. In particular, the subapplication can be used as the CmdType of the application, and/or term frequency analysis (e.g., TF/IDF) can be used on command line information to group, for example, any marathon related applications together (e.g., as a python.marathon CmdType) and separately from other Python applications (e.g., as a python.airflow CmdType).

In various embodiments, machine learning techniques are used to determine a CmdType. The CmdType model is constrained such that the execution path for each CmdType is unique. One example approach to making a CmdType model is a random forest based approach. An initial CmdType model is bootstrapped using process parameters (e.g., available within one minute of process startup) obtained using one hour of information for a given customer (e.g., entity A). Examples of such parameters include the command line of the process, the command line of the process’s parent(s) (if applicable), the uptime of the process, UID/EUID and any change information, TTY and any change information, listening ports, and children (if any). Another approach is to perform term frequency clustering over command line information to convert command lines into cluster identifiers.

The random forest model can be used (e.g., in subsequent hours) to predict a CmdType for a process (e.g., based on features of the process). If a match is found, the process can be assigned the matching CmdType. If a match is not found, a comparison between features of the process and its nearest CmdType (e.g., as determined using a Levenstein distance) can be performed. The existing CmdType can be expanded to include the process, or, as applicable, a new CmdType can be created (and other actions taken, such as generating an alert). Another approach to handling processes which do not match an existing CmdType is to designate such processes as unclassified, and once an hour, create a new random forest seeded with process information from a sampling of classified processes (e.g., 10 or 100 processes per CmdType) and the new processes. If a given new process winds up in an existing set, the process is given the corresponding CmdType. If a new cluster is created, a new CmdType can be created.

Conceptually, a polygraph represents the smallest possible graph of clusters that preserve a set of rules (e.g., in which nodes included in the cluster must share a CmdType and behavior). As a result of performing MNC, SimRank, and cluster splitting (e.g., CmdType splitting) many processes are clustered together based on commonality of behavior (e.g., communication behavior) and commonality of application type. Such clustering represents a significant reduction in graph size (e.g., compared to the original raw physical graph). Nonetheless, further clustering can be performed (e.g., by iterating on the graph data using the GBM to achieve such a polygraph). As more information within the graph is correlated, more nodes can be clustered together, reducing the size of the graph, until convergence is reached and no further clustering is possible.

FIG. 3E depicts two pairs of clusters. In particular, cluster 320 represents a set of client processes sharing the same CmdType (“a1”), communicating (collectively) with a server process having a CmdType (“a2”). Cluster 322 also represents a set of client processes having a CmdType a1 communicating with a server process having a CmdType a2.

33

The nodes in clusters **320** and **322** (and similarly nodes in **321** and **323**) remain separately clustered (as depicted) after MNC/SimRank/CmdType splitting-isolated islands. One reason this could occur is where server process **321** corresponds to processes executing on a first machine (having an IP address of 1.1.1.1). The machine fails and a new server process **323** starts, on a second machine (having an IP address of 2.2.2.2) and takes over for process **321**.

Communications between a cluster of nodes (e.g., nodes of cluster **320**) and the first IP address can be considered different behavior from communications between the same set of nodes and the second IP address, and thus communications **324** and **325** will not be combined by MNC/SimRank in various embodiments. Nonetheless, it could be desirable for nodes of clusters **320/322** to be combined (into cluster **326**), and for nodes of clusters **321/323** to be combined (into cluster **327**), as representing (collectively) communications between a1 and a2. One task that can be performed by data platform **12** is to use DNS query information to map IP addresses to logical entities. As will be described in more detail below, GBM **154** can make use of the DNS query information to determine that graph nodes of cluster **320** and graph nodes of cluster **322** both made DNS queries for “appserverabc.example.com,” which first resolved to 1.1.1.1 and then to 2.2.2.2, and to combine nodes **320/322** and **321/323** together into a single pair of nodes (**326** communicating with **327**).

In various embodiments, GBM **154** operates in a batch manner in which it receives as input the nodes and edges of a graph for a particular time period along with its previous state, and generates as output clustered nodes, cluster membership edges, cluster-to-cluster edges, events, and its next state.

GBM **154** may not try to consider all types of entities and their relationships that may be available in a conceptual common graph all at once. Instead, GBM uses a concept of models where a subset of node and edge types and their properties are considered in a given model. Such an approach is helpful for scalability, and also to help preserve detailed information (of particular importance in a security context)—as clustering entities in a more complex and larger graph could result in less useful results. In particular, such an approach allows for different types of relationships between entities to be preserved/more easily analyzed.

While GBM **154** can be used with different models corresponding to different subgraphs, core abstractions remain the same across types of models.

For example, each node type in a GBM model is considered to belong to a class. The class can be thought of as a way for the GBM to split nodes based on the criteria it uses for the model. The class for a node is represented as a string whose value is derived from the node’s key and properties depending on the GBM Model. Note that different GBM models may create different class values for the same node. For each node type in a given GBM model, GBM **154** can generate clusters of nodes for that type. A GBM generated cluster for a given member node type cannot span more than one class for that node type. GBM **154** generates edges between clusters that have the same types as the edges between source and destination cluster node types.

Additionally or alternatively, the processes described herein as being used for a particular model can be used (can be the same) across models, and different models can also be configured with different settings.

Additionally or alternatively, the node types and the edge types may correspond to existing types in the common graph node and edge tables but this is not necessary. Even when

34

there is a correspondence, the properties provided to GBM **154** are not limited to the properties that are stored in the corresponding graph table entries. They can be enriched with additional information before being passed to GBM **154**.

Logically, the input for a GBM model can be characterized in a manner that is similar to other graphs. Edge triplets can be expressed, for example, as an array of source node type, edge type, and destination node type. And, each node type is associated with node properties, and each edge type is associated with edge properties. Other edge triplets can also be used (and/or edge triplets can be extended) in accordance with various embodiments.

Note that the physical input to the GBM model need not (and does not, in various embodiments) conform to the logical input. For example, the edges in the PtypeConn model correspond to edges between Matching Neighbors (MN) clusters, where each process node has an MN cluster identifier property. In the User ID to User ID model (also referred to herein as the Uid2Uid model), edges are not explicitly provided separately from nodes (as the euid array in the node properties serves the same purpose). In both cases, however, the physical information provides the applicable information necessary for the logical input.

The state input for a particular GBM model can be stored in a file, a database, or other appropriate storage. The state file (from a previous run) is provided, along with graph data, except for when the first run for a given model is performed, or the model is reset. In some cases, no data may be available for a particular model in a given time period, and GBM may not be run for that time period. As data becomes available at a future time, GBM can run using the latest state file as input.

GBM **154** outputs cluster nodes, cluster membership edges, and inter-cluster relationship edges that are stored (in some embodiments) in the graph node tables: node_c, node_cm, and node_icr, respectively. The type names of nodes and edges may conform to the following rules:

A given node type can be used in multiple different GBM models. The type names of the cluster nodes generated by two such models for that node type will be different. For instance, process type nodes will appear in both PtypeConn and Uid2Uid models, but their cluster nodes will have different type names.

The membership edge type name is “MemberOf.”

The edge type names for cluster-to-cluster edges will be the same as the edge type names in the underlying node-to-node edges in the input.

The following are example events GBM **154** can generate: new class, new cluster, new edge from class to class, split class (the notion that GBM **154** considers all nodes of a given type and class to be in the same cluster initially and if GBM **154** splits them into multiple clusters, it is splitting a class), new edge from cluster and class, new edge between cluster and cluster, and/or new edge from class to cluster.

One underlying node or edge in the logical input can cause multiple types of events to be generated. Conversely, one event can correspond to multiple nodes or edges in the input. Not every model generates every event type.

Additional information regarding examples of data structures/models that can be used in conjunction with models used by data platform **12** is now provided.

In some examples, a PTypeConn model clusters nodes of the same class that have similar connectivity relationships. For example, if two processes had similar incoming neighbors of the same class and outgoing neighbors of the same class, they could be clustered.

35

The node input to the PTypeConn model for a given time period includes non-interactive (i.e., not associated with tty) process nodes that had connections in the time period and the base graph nodes of other types (IP Service Endpoint (IPSep) comprising an IP address and a port, DNS Service Endpoint (DNSSep) and IP Address) that have been involved in those connections.

The base relationship is the connectivity relationship for the following type triplets:

Process, ConnectedTo, Process
 Process, ConnectedTo, IP Service Endpoint (IPSep)
 Process, ConnectedTo, DNS Service Endpoint (DNSSep)
 IP Address, ConnectedTo, ProcessProcess, DNS, ConnectedTo, Process

The edge inputs to this model are the ConnectedTo edges from the MN cluster, instead of individual node-to-node ConnectedTo edges from the base graph. The membership edges created by this model refer to the base graph node type provided in the input.

Class Values:

The class values of nodes are determined as follows depending on the node type (e.g., Process nodes, IPSep nodes, DNSSep nodes, and IP Address nodes).

Process Nodes:

if exe_path contains java then "java <cmdline term 1> . . ."
 else if exe_path contains python then "python <cmdline term 1> . . ."
 else "last_part_of_exe_path"

IPSep Nodes:

if IP_internal then "IntIPS"
 else if severity=0 then "<IP_addr>:<protocol>:<port>"
 else "<IP_addr>:<port>BadIP"

DNSSep Nodes:

if IP_internal=1 then "<hostname>"
 else if severity=0 then "<hostname>:<protocol>:<port>"
 else "<hostname>:<port>_BadIP"

IP Address Nodes (Will Appear Only on Client Side):

if IP_internal=1 then "IPIntC"
 else if severity=0 then "ExtIPC"
 else "ExtBadIPC"

Events:

A new class event in this model for a process node is equivalent to seeing a new CType being involved in a connection for the first time. Note that this does not mean the CType was not seen before. It is possible that it was previously seen but did not make a connection at that time.

A new class event in this model for an IPSep node with IP_internal=0 is equivalent to seeing a connection to a new external IP address for the first time.

A new class event in this model for a DNSSep node is equivalent to seeing a connection to a new domain for the first time.

A new class event in this model for an IP Address node with IP_internal=0 and severity=0 is equivalent to seeing a connection from any external IP address for the first time.

A new class event in this model for an IP Address node with IP internal=0 and severity>0 is equivalent to seeing a connection from any bad external IP address for the first time.

A new class to class edge from a class for a process node to a class for a process node is equivalent to seeing a communication from the source CType making a connection to the destination CType for the first time.

A new class to class edge from a class for a process node to a class for a DNSSep node is equivalent to seeing a

36

communication from the source CType making a connection to the destination domain name for the first time.

An IntPConn Model may be similar to the PtypeConn Model, except that connection edges between parent/child processes and connections between processes where both sides are not interactive are filtered out.

A Uid2Uid Model may cluster processes with the same username that show similar privilege change behavior. For instance, if two processes with the same username had similar effective user values, launched processes with similar usernames, and were launched by processes with similar usernames, then they could be clustered.

An edge between a source cluster and destination cluster generated by this model means that all of the processes in the source cluster had a privilege change relationship to at least one process in the destination cluster.

The node input to this model for a given time period includes process nodes that are running in that period. The value of a class of process nodes is "<username>".

The base relationship that is used for clustering is privilege change, either by the process changing its effective user ID, or by launching a child process which runs with a different user.

The physical input for this model includes process nodes (only), with the caveat that the complete ancestor hierarchy of process nodes active (i.e., running) for a given time period is provided as input even if an ancestor is not active in that time period. Note that effective user IDs of a process are represented as an array in the process node properties, and launch relationships are available from ppid_hash fields in the properties as well.

A new class event in this model is equivalent to seeing a user for the first time.

A new class to class edge event is equivalent to seeing the source user making a privilege change to the destination user for the first time.

A Ct2Ct Model may cluster processes with the same CType that show similar launch behavior. For instance, if two processes with the same CType have launched processes with similar CTypes, then they could be clustered.

The node input to this model for a given time period includes process nodes that are running in that period. The value class of process nodes is CType (similar to how it is created for the PtypeConn Model).

The base relationship that is used for clustering is a parent process with a given CType launching a child process with another given destination CType.

The physical input for this model includes process nodes (only) with the caveat that the complete ancestor hierarchy active process nodes (i.e., that are running) for a given time period is provided as input even if an ancestor is not active in that time period. Note that launch relationships are available from ppid_hash fields in the process node properties.

An edge between a source cluster and destination cluster generated by this model means that all of the processes in the source cluster launched at least one process in the destination cluster.

A new class event in this model is equivalent to seeing a CType for the first time. Note that the same type of event will be generated by the PtypeConn Model as well.

A new class to class edge event is equivalent to seeing the source CType launching the destination CType for the first time.

An MTypeConn Model may cluster nodes of the same class that have similar connectivity relationships. For

example, if two machines had similar incoming neighbors of the same class and outgoing neighbors of the same class, they could be clustered.

A new class event in this model will be generated for external IP addresses or (as applicable) domain names seen for the first time. Note that a new class to class to edge Machine, class to class for an IP Sep or DNSName node will also be generated at the same time.

The membership edges generated by this model will refer to Machine, IP Address, DNSName, and IP Sep nodes in the base graph. Though the nodes provided to this model are IP Address nodes instead of IP Sep nodes, the membership edges it generates will refer to IP Sep type nodes. Alternatively, the base graph can generate edges between Machine and IP Sep node types. Note that the Machine to IP Address edges have tcp_dst_ports/udp_dst_ports properties that can be used for this purpose.

The node input to this model for a given time period includes machine nodes that had connections in the time period and the base graph nodes of other types (IP Address and DNSName) that were involved in those connections.

The base relationship is the connectivity relationship for the following type triplets:

Machine, ConnectedTo, Machine
Machine, ConnectedTo, IP Address
Machine, ConnectedTo, DNSName
IP Address, ConnectedTo, Machine, DNS, ConnectedTo, Machine

The edge inputs to this model are the corresponding ConnectedTo edges in the base graph.

Class Values:

Machine:

The class value for all Machine nodes is "Machine."

The machine_terms property in the Machine nodes is used, in various embodiments, for labeling machines that are clustered together. If a majority of the machines clustered together share a term in the machine_terms, that term can be used for labeling the cluster.

IP Sep:

The class value for IP Sep nodes is determined as follows:
if IP_internal then "IntIPS"
else
if severity=0 then "<ip_addr>:<protocol>:<port>"
else "<IP_addr_BadIP>"

IP Address:

The class value for IP Address nodes is determined as follows:

if IP_internal then "IntIPC"
else
if severity=0 then "ExtIPC"
else "ExtBadIPC"

DNSName:

The class value for DNSName nodes is determined as follows:

if severity=0 then "<hostname>"
else then "<hostname>_BadIP"

An example structure for a New Class Event is now described.

The key field for this event type looks as follows (using the PtypeConn model as an example):

```
{
  "node": {
    "class": {
      "cid": "httpd"
    },

```

```
    "key": {
      "cid": "29654"
    },
    "type": "PtypeConn"
  }
}
```

It contains the class value and also the ID of the cluster where that class value is observed. Multiple clusters can be observed with the same value in a given time period. It contains the class value and also the ID of the cluster where that class value is observed. Multiple clusters can be observed with the same value in a given time period. Accordingly, in some embodiments, GBM 154 generates multiple events of this type for the same class value.

The properties field looks as follows:

```
{
  "set_size": 5
}
```

The set_size indicates the size of the cluster referenced in the keys field.

Conditions:

For a given model and time period, multiple NewClass events can be generated if there is more than one cluster in that class. NewNode events will not be generated separately in this case. Example New Class to Class Edge Event structure:

The key field for this event type looks as follows (using the PtypeConn model as an example):

```
"edge": {
  "dst_node": {
    "class": {
      "cid": "java war"
    },
    "key": {
      "cid": "27635"
    },
    "type": "PtypeConn"
  },
  "src_node": {
    "class": {
      "cid": "IntIPC"
    },
    "key": {
      "cid": "20881"
    },
    "type": "PtypeConn"
  },
  "type": "ConnectedTo"
}
```

The key field contains source and destination class values and also source and destination cluster identifiers (i.e., the src/dst_node: key.cid represents the src/dst cluster identifier).

In a given time period for a given model, an event of this type could involve multiple edges between different cluster pairs that have the same source and destination class values. GBM 154 can generate multiple events in this case with different source and destination cluster identifiers.

The props fields look as follows for this event type:

```
{
  "dst_set_size": 2,
  "src_set_size": 1
}
```

The source and destination sizes represent the sizes of the clusters given in the keys field.

Conditions:

For a given model and time period, multiple NewClassTo-Class events can be generated if there are more than one pair

of clusters in that class pair. NewNodeToNode events are not generated separately in this case.

Combining Events at the Class Level: for a given model and time period, the following example types of events can represent multiple changes in the underlying GBM cluster level graph in terms of multiple new clusters or multiple new edges between clusters:

NewClass
NewEdgeClassToClass
NewEdgeNodeToClass
NewEdgeClassToNode

Multiple NewClass events with the same model and class can be output if there are multiple clusters in that new class.

Multiple NewEdgeClassToClass events with the same model and class pair can be output if there are multiple new cluster edges within that class pair.

Multiple NewEdgeNodeToClass events with the same model and destination class can be output if there are multiple new edges from the source cluster to the destination clusters in that destination class (the first time seeing this class as a destination cluster class for the source cluster).

Multiple NewEdgeClassToNode events with the same model and source class can be output if there are multiple new edges from source clusters to the destination clusters in that source class (the first time seeing this class as a source cluster class for the destination cluster).

These events may be combined at the class level and treated as a single event when it is desirable to view changes at the class level, e.g., when one wants to know when there is a new CType.

In some examples, different models may have partial overlap in the types of nodes they use from the base graph. Therefore, they can generate NewClass type events for the same class. NewClass events can also be combined across models when it is desirable to view changes at the class level.

Using techniques herein, actions can be associated with processes and (e.g., by associating processes with users) actions can thus also be associated with extended user sessions. Such information can be used to track user behavior correctly, even where a malicious user attempts to hide his trail by changing user identities (e.g., through lateral movement). Extended user session tracking can also be useful in operational use cases without malicious intent, e.g., where users make original logins with distinct usernames (e.g., "charlie" or "dave") but then perform actions under a common username (e.g., "admin" or "support"). One such example is where multiple users with administrator privileges exist, and they need to gain superuser privilege to perform a particular type of maintenance. It may be desirable to know which operations are performed (as the superuser) by which original user when debugging issues. In the following examples describing extended user session tracking, reference is generally made to using the secure shell (ssh) protocol as implemented by openssh (on the server side) as the mechanism for logins. However, extended user session tracking is not limited to the ssh protocol or a particular limitation and the techniques described herein can be extended to other login mechanisms.

On any given machine, there will be a process that listens for and accepts ssh connections on a given port. This process can run the openssh server program running in daemon mode or it could be running another program (e.g., initd on a Linux system). In either case, a new process running openssh will be created for every new ssh login session and

this process can be used to identify an ssh session on that machine. This process is called the "privileged" process in openssh.

After authentication of the ssh session, when an ssh client requests a shell or any other program to be run under that ssh session, a new process that runs that program will be created under (i.e., as a child of) the associated privileged process. If an ssh client requests port forwarding to be performed, the connections will be associated with the privileged process.

In modern operating systems such as Linux and Windows, each process has a parent process (except for the very first process) and when a new process is created the parent process is known. By tracking the parent-child hierarchy of processes, one can determine if a particular process is a descendant of a privileged openssh process and thus if it is associated with an ssh login session.

For user session tracking across machines (or on a single machine with multiple logins) in a distributed environment, it is established when two login sessions have a parent-child relationship. After that, the "original" login session, if any, for any given login session can be determined by following the parent relationship recursively.

FIG. 3F is a representation of a user logging into a first machine and then into a second machine from the first machine, as well as information associated with such actions. In the example of FIG. 3F, a user, Charlie, logs into Machine A (331) from a first IP address (332). As part of the login process, he provides a username (333). Once connected to Machine A, an openssh privileged process (334) is created to handle the connection for the user, and a terminal session is created and a bash process (335) is created as a child. Charlie launches an ssh client (336) from the shell, and uses it to connect (337) to Machine B (338). As with the connection he makes to Machine A, Charlie's connection to Machine B will have an associated incoming IP address (339), in this case, the IP address of Machine A. And, as part of the login process with Machine B, Charlie will provide a username (340) which need not be the same as username 333. An openssh privileged process (341) is created to handle the connection, and a terminal session and child bash process (342) will be created. From the command line of Machine B, Charlie launches a curl command (343), which opens an HTTP connection (344) to an external Machine C (345).

FIG. 3G is an alternate representation of actions occurring in FIG. 3F, where events occurring on Machine A are indicated along line 350, and events occurring on Machine B are indicated along line 351. As shown in FIG. 3G, an incoming ssh connection is received at Machine A (352). Charlie logs in (as user "x") and an ssh privileged process is created to handle Charlie's connection (353). A terminal session is created and a bash process is created (354) as a child of process 353. Charlie wants to ssh to Machine B, and so executes an ssh client on Machine A (355), providing credentials (as user "y") at 356. Charlie logs into Machine B, and an ssh privileged process is created to handle Charlie's connection (357). A terminal session is created and a bash process is created (358) as a child of process 357. Charlie then executes curl (359) to download content from an external domain (via connection 360).

The external domain could be a malicious domain, or it could be benign. Suppose the external domain is malicious (and, e.g., Charlie has malicious intent). It would be advantageous (e.g., for security reasons) to be able to trace the contact with the external domain back to Machine A, and then back to Charlie's IP address. Using techniques described herein (e.g., by correlating process information

41

collected by various agents), such tracking of Charlie's activities back to his original login (330) can be accomplished. In particular, an extended user session can be tracked that associates Charlie's ssh processes together with a single original login and thus original user.

As described herein, software agents (such as agent 112) may run on machines (such as a machine that implements one of nodes 116) and detect new connections, processes, and/or logins. As also previously explained, such agents send associated records to data platform 12 which includes one or more datastores (e.g., data store 30) for persistently storing such data. Such data can be modeled using logical tables, also persisted in datastores (e.g., in a relational database that provides an SQL interface), allowing for querying of the data. Other datastores such as graph oriented databases and/or hybrid schemes can also be used.

The following identifiers are commonly used in the tables:

MID

PID_hash

An ssh login session can be identified uniquely by an (MID, PID_hash) tuple. The MID is a machine identifier that is unique to each machine, whether physical or virtual, across time and space. Operating systems use numbers called process identifiers (PIDs) to identify processes running at a given time. Over time processes may die and new processes may be started on a machine or the machine itself may restart. The PID is not necessarily unique across time in that the same PID value can be reused for different processes at different times. In order to track process descendants across time, one should therefore account for time as well. In order to be able to identify a process on a machine uniquely across time, another number called a PID hash is generated for the process. In various embodiments, the PID_hash is generated using a collision-resistant hash function that takes the PID, start time, and (in various embodiments, as applicable) other properties of a process.

Input data collected by agents comprises the input data model and is represented by the following logical tables:

connections

processes

logins

A connections table may maintain records of TCP/IP connections observed on each machine. Example columns included in a connections table are as follows:

Column Name	Description
MID	Identifier of the machine that the connection was observed on.
start_time	Connection start time.
PID_hash	Identifier of the process that was associated with the connection.
src_IP_addr	Source IP address (the connection was initiated from this IP address).
src_port	Source port.
dst_IP_addr	Destination IP address (the connection was made to this IP address).
dst_port	Destination port.
Prot	Protocol (TCP or UDP).
Dir	Direction of the connection (incoming or outgoing) with respect to this machine.

The source fields (IP address and port) correspond to the side from which the connection was initiated. On the destination side, the agent associates an ssh connection with the privileged ssh process that is created for that connection.

For each connection in the system, there will be two records in the table, assuming that the machines on both

42

sides of the connection capture the connection. These records can be matched based on equality of the tuple (src_IP_addr, src_port, dst_IP_addr, dst_port, Prot) and proximity of the start_time fields (e.g., with a one minute upper threshold between the start time fields).

A processes table maintains records of processes observed on each machine. It may have the following columns:

Column Name	Description
MID	Identifier of the machine that the process was observed on.
PID_hash	Identifier of the process.
start_time	Start time of the process.
exe_path	The executable path of the process.
PPID_hash	Identifier of the parent process.

A logins table may maintain records of logins to machines. It may have the following columns:

Column Name	Description
MID	Identifier of the machine that the login was observed on.
sshd_PID_hash	Identifier of the sshd privileged process associated with login.
login_time	Time of login.
login_username	Username used in login.

Output data generated by session tracking is represented with the following logical tables:

login-local-descendant

login-connection

login-lineage

Using data in these tables, it is possible to determine descendant processes of a given ssh login session across the environment (i.e., spanning machines). Conversely, given a process, it is possible to determine if it is an ssh login descendant as well as the original ssh login session for it if so.

A login-local-descendant table maintains the local (i.e., on the same machine) descendant processes of each ssh login session. It may have the following columns:

Column Name	Description
MID	Identifier of the machine that the login was observed on.
sshd_PID_hash	Identifier of the sshd privileged process associated with login.
login_time	Time of login.
login_username	Username used in login.

A login-connections table may maintain the connections associated with ssh logins. It may have the following columns:

Column Name	Description
MID	Identifier of the machine that the process was observed on.
sshd_PID_hash	Identifier of the sshd privileged process associated with the login.
login_time	Time of login.
login_username	The username used in the login.
src_IP_addr	Source IP address (connection was initiated from this IP address).

-continued

Column Name	Description
src_port	Source port.
dst_IP_addr	Destination IP address (connection was made to this IP address).
dst_port	Destination port.

A login-lineage table may maintain the lineage of ssh login sessions. It may have the following columns:

Column Name	Description
MID	Identifier of the machine that the ssh login was observed on.
sshd_PID_hash	Identifier of the sshd privileged process associated with the login.
parent_MID	Identifier of the machine that the parent ssh login was observed on.
parent_sshd_PID_hash	Identifier of the sshd privileged process associated with the parent login.
origin_MID	Identifier of the machine that the origin ssh login was observed on.
origin_sshd_PID_hash	Identifier of the sshd privileged process associated with the origin login.

The parent_MID and parent_sshd_PID_hash columns can be null if there is no parent ssh login. In that case, the (MID, sshd_PID_hash) tuple will be the same as the (origin_MID, origin_sshd_PID_hash) tuple.

FIG. 3H illustrates an example of a process for performing extended user tracking. In various embodiments, process 361 is performed by data platform 12. The process begins at 362 when data associated with activities occurring in a network environment (such as entity A's datacenter) is received. One example of such data that can be received at 362 is agent-collected data described above (e.g., in conjunction with process 200). At 363, the received network activity is used to identify user login activity. And, at 364, a logical graph that links the user login activity to at least one user and at least one process is generated (or updated, as applicable). Additional detail regarding process 361, and in particular, portions 363 and 364 of process 361 are described in more detail below (e.g., in conjunction with discussion of FIG. 3J).

FIG. 3I depicts a representation of a user logging into a first machine, then into a second machine from the first machine, and then making an external connection. The scenario depicted in FIG. 3I is used to describe an example of processing that can be performed on data collected by agents to generate extended user session tracking information. FIG. 3I is an alternate depiction of the information shown in FIGS. 3F and 3G.

At time t1 (365), a first ssh connection is made to Machine A (366) from an external source (367) by a user having a username of "X." In the following example, suppose the external source has an IP address of 1.1.1.10 and uses source port 10000 to connect to Machine A (which has an IP address of 2.2.2.20 and a destination port 22). External source 367 is considered an external source because its IP address is outside of the environment being monitored (e.g., is a node outside of entity A's datacenter, connecting to a node inside of entity A's datacenter).

A first ssh login session LS1 is created on machine A for user X. The privileged openssh process for this login is A1 (368). Under the login session LS1, the user creates a bash shell process with PID_hash A2 (369).

At time t2 (370), inside the bash shell process A2, the user runs an ssh program under a new process A3 (371) to log in to machine B (372) with a different username ("Y"). In particular, an ssh connection is made from source IP address 2.2.2.20 and source port 10001 (Machine A's source information) to destination IP address 2.2.2.21 and destination port 22 (Machine B's destination information).

A second ssh login session LS2 is created on machine B for user Y. The privileged openssh process for this login is B1 (373). Under the login session LS2, the user creates a bash shell process with PID_hash B2 (374).

At time t3 (376), inside the bash shell process B2, the user runs a curl command under a new process B3 (377) to download a file from an external destination (378). In particular, an HTTPS connection is made from source IP address 2.2.2.21 and source port 10002 (Machine B's source information) to external destination IP address 3.3.3.30 and destination port 443 (the external destination's information).

Using techniques described herein, it is possible to determine the original user who initiated the connection to external destination 378, which in this example is a user having the username X on machine A (where the extended user session can be determined to start with ssh login session LS1).

Based on local descendant tracking, the following determinations can be on machine A and B without yet having performed additional processing (described in more detail below):

A3 is a descendant of A1 and thus associated with LS1.

The connection to the external domain from machine B is initiated by B3.

B3 is a descendant of B1 and is thus associated with LS2. Connection to the external domain is thus associated with LS2.

An association between A3 and LS2 can be established based on the fact that LS2 was created based on an ssh connection initiated from A3. Accordingly, it can be determined that LS2 is a child of LS1.

To determine the user responsible for making the connection to the external destination (e.g., if it were a known bad destination), first, the process that made the connection would be traced, i.e., from B3 to LS2. Then LS2 would be traced to LS1 (i.e., LS1 is the origin login session for LS2). Thus the user for this connection is the user for LS1, i.e., X. As represented in FIG. 3I, one can visualize the tracing by following the links (in the reverse direction of arrows) from external destination 378 to A1 (368).

In the example scenario, it is assumed that both ssh connections occur in the same analysis period. However, the approaches described herein will also work for connections and processes that are created in different time periods.

FIG. 3J illustrates an example of a process for performing extended user tracking. In various embodiments, process 380 is performed periodically (e.g., once an hour in a batch fashion) by ssh tracker 148 to generate new output data. In general, batch processing allows for efficient analysis of large volumes of data. However, the approach can be adapted, as applicable, to process input data on a record-by-record fashion while maintaining the same logical data processing flow. As applicable the results of a given portion of process 380 are stored for use in a subsequent portion.

The process begins at 381 when new ssh connection records are identified. In particular, new ssh connections started during the current time period are identified by querying the connections table. The query uses filters on the start_time and dst_port columns. The values of the range filter on the start_time column are based on the current time

period. The `dst_port` column is checked against `ssh` listening port(s). By default, the `ssh` listening port number is 22. However, as this could vary across environments, the port(s) that openssh servers are listening to in the environment can be determined by data collection agents dynamically and used as the filter value for the `dst` port as applicable. In the scenario depicted in FIG. 3I, the query result will generate the records shown in FIG. 3K. Note that for the connection between machine A and B, the two machines are likely to report `start_time` values that are not exactly the same but close enough to be considered matching (e.g., within one minute or another appropriate amount of time). In the above table, they are shown to be the same for simplicity.

At **382**, `ssh` connection records reported from source and destination sides of the same connection are matched. The `ssh` connection records (e.g., returned from the query at **381**) are matched based on the following criteria:

The five tuples (`src_IP`, `dst_IP`, `IP_prot`, `src_port`, `dst_port`) of the connection records must match.

The delta between the start times of the connections must be within a limit that would account for the worst case clock difference expected between two machines in the environment and typical connection setup latency.

If there are multiple matches possible, then the match with the smallest time delta is chosen.

Note that record **390** from machine A for the incoming connection from the external source cannot be matched with another record as there is an agent only on the destination side for this connection. Example output of portion **382** of process **380** is shown in FIG. 3L. The values in the `dst_PID_hash` column (**391**) are that of the `sshd` privileged process associated with `ssh` logins.

At **383**, new logins during the current time period are identified by querying the logins table. The query uses a range filter on the `login_time` column with values based on the current time period. In the example depicted in FIG. 3I, the query result will generate the records depicted in FIG. 3M.

At **384**, matched `ssh` connection records created at **382** and new login records created at **383** are joined to create new records that will eventually be stored in the login-connection table. The join condition is that `dst MID` of the matched connection record is equal to the `MID` of the login record and the `dst PID hash` of the matched connection record is equal to the `sshd PID_hash` of the login record. In the example depicted in FIG. 3I, the processing performed at **384** will generate the records depicted in FIG. 3N.

At **385**, login-local-descendant records in the lookback time period are identified. It is possible that a process that is created in a previous time period makes an `ssh` connection in the current analysis batch period. Although not depicted in the example illustrated in FIG. 3I, consider a case where `bash` process A2 does not create `ssh` process A3 right away but instead that the `ssh` connection A3 later makes to machine B is processed in a subsequent time period than the one where A2 was processed. While processing this subsequent time period in which processes A3 and B1 are seen, knowledge of A2 would be useful in establishing that B1 is associated with A3 (via `ssh` connection) which is associated with A2 (via process parentage) which in turn would be useful in establishing that the parent of the second `ssh` login is the first `ssh` login. The time period for which look back is performed can be limited to reduce the amount of historical data that is considered. However, this is not a requirement (and the amount of look back can be determined, e.g., based on available processing resources). The login local descendants in the lookback time period can be identified by

querying the login-local-descendant table. The query uses a range filter on the `login_time` column where the range is from `start_time_of_current_period-lookback_time` to `start_time_of_current_period`. (No records as a result of performing **385** on the scenario depicted in FIG. 3I are obtained, as only a single time period is applicable in the example scenario.)

At **386**, new processes that are started in the current time period are identified by querying the processes table. The query uses a range filter on the `start_time` column with values based on the current time period. In the example depicted in FIG. 3I, the processing performed at **386** will generate the records depicted in FIG. 3O.

At **387**, new login-local-descendant records are identified. The purpose is to determine whether any of the new processes in the current time period are descendants of an `ssh` login process and if so to create records that will be stored in the login-local-descendant table for them. In order to do so, the parent-child relationships between the processes are recursively followed. Either a top down or bottom up approach can be used. In a top down approach, the `ssh` local descendants in the lookback period identified at **385**, along with new `ssh` login processes in the current period identified at **384** are considered as possible ancestors for the new processes in the current period identified at **386**.

Conceptually, the recursive approach can be considered to include multiple sub-steps where new processes that are identified to be `ssh` local descendants in the current sub-step are considered as ancestors for the next step. In the example scenario depicted in FIG. 3I, the following descendancy relationships will be established in two sub-steps:

Sub-Step 1:

Process A2 is a local descendant of LS1 (i.e., `MID`=A, `sshd_PID_hash`=A1) because it is a child of process A1 which is the login process for LS1.

Process B2 is a local descendant of LS2 (i.e., `MID`=B, `sshd_PID_hash`=B1) because it is a child of process B1 which is the login process for LS2.

Sub-step 2:

Process A3 is a local descendant of LS1 because it is a child of process A2 which is associated to LS1 in sub-step 1.

Process B3 is a local descendant of LS2 because it is a child of process B1 which is associated to LS2 in sub-step 1.

Implementation portion **387** can use a datastore that supports recursive query capabilities, or, queries can be constructed to process multiple conceptual sub-steps at once. In the example depicted in FIG. 3I, the processing performed at **387** will generate the records depicted in FIG. 3P. Note that the `ssh` privileged processes associated with the logins are also included as they are part of the login session.

At **388**, the lineage of new `ssh` logins created in the current time period is determined by associating their `ssh` connections to source processes that may be descendants of other `ssh` logins (which may have been created in the current period or previous time periods). In order to do so, first an attempt is made to join the new `ssh` login connections in the current period (identified at **384**) with the combination of the login local descendants in the lookback period (identified at **385**) and the login local descendants in the current time period (identified at **386**). This will create adjacency relationships between child and parent logins. In the example depicted in FIG. 3I, the second `ssh` login connection will be associated with process A3 and an adjacency relationship between the two login sessions will be created (as illustrated in FIG. 3Q).

Next, the adjacency relationships are used to find the original login sessions. While not shown in the sample scenario, there could be multiple ssh logins in a chain in the current time period, in which case a recursive approach (as in 387) could be used. At the conclusion of portion 388, the login lineage records depicted in FIG. 3R will be generated.

Finally, at 389, output data is generated. In particular, the new login-connection, login-local-descendant, and login-lineage records generated at 384, 387, and 388 are inserted into their respective output tables (e.g., in a transaction manner).

An alternate approach to matching TCP connections between machines running an agent is for the client to generate a connection GUID and send it in the connection request (e.g., the SYN packet) it sends and for the server to extract the GUID from the request. If two connection records from two machines have the same GUID, they are for the same connection. Both the client and server will store the GUID (if it exists) in the connection records they maintain and report. On the client side, the agent can configure the network stack (e.g. using IP tables functionality on Linux) to intercept an outgoing TCP SYN packet and modify it to add the generated GUID as a TCP option. On the server side, the agent already extracts TCP SYN packets and thus can look for this option and extract the GUID if it exists.

Example graph-based user tracking and threat detection embodiments associated with data platform 12 will now be described. Administrators and other users of network environments (e.g., entity A's datacenter 104) often change roles to perform tasks. As one example, suppose that at the start of a workday, an administrator (hereinafter "Joe Smith") logs in to a console, using an individualized account (e.g., username=joe.smith). Joe performs various tasks as himself (e.g., answering emails, generating status reports, writing code, etc.). For other tasks (e.g., performing updates), Joe may require different/additional permission than his individual account has (e.g., root privileges). One way Joe can gain access to such permissions is by using sudo, which will allow Joe to run a single command with root privileges. Another way Joe can gain access to such permissions is by su or otherwise logging into a shell as root. After gaining root privileges, another thing that Joe can do is switch identities. As one example, to perform administrative tasks, Joe may use "su help" or "su database-admin" to become (respectively) the help user or the database-admin user on a system. He may also connect from one machine to another, potentially changing identities along the way (e.g., logging in as joe.smith at a first console, and connecting to a database server as database-admin). When he's completed various administrative tasks, Joe can relinquish his root privileges by closing out of any additional shells created, reverting back to a shell created for user joe.smith.

While there are many legitimate reasons for Joe to change his identity throughout the day, such changes may also correspond to nefarious activity. Joe himself may be nefarious, or Joe's account (joe.smith) may have been compromised by a third party (whether an "outsider" outside of entity A's network, or an "insider"). Using techniques described herein, the behavior of users of the environment can be tracked (including across multiple accounts and/or multiple machines) and modeled (e.g., using various graphs described herein). Such models can be used to generate alerts (e.g., to anomalous user behavior). Such models can also be used forensically, e.g., helping an investigator visualize various aspects of a network and activities that have

occurred, and to attribute particular types of actions (e.g., network connections or file accesses) to specific users.

In a typical day in a datacenter, a user (e.g., Joe Smith) will log in, run various processes, and (optionally) log out. The user will typically log in from the same set of IP addresses, from IP addresses within the same geographical area (e.g., city or country), or from historically known IP addresses/geographical areas (i.e., ones the user has previously/occasionally used). A deviation from the user's typical (or historical) behavior indicates a change in login behavior. However, it does not necessarily mean that a breach has occurred. Once logged into a datacenter, a user may take a variety of actions. As a first example, a user might execute a binary/script. Such binary/script might communicate with other nodes in the datacenter, or outside of the datacenter, and transfer data to the user (e.g., executing "curl" to obtain data from a service external to the datacenter). As a second example, the user can similarly transfer data (e.g., out of the datacenter), such as by using POST. As a third example, a user might change privilege (one or more times), at which point the user can send/receive data as per above. As a fourth example, a user might connect to a different machine within the datacenter (one or more times), at which point the user can send/receive data as per the above.

In various embodiments, the above information associated with user behavior is broken into four tiers. The tiers represent example types of information that data platform 12 can use in modeling user behavior:

1. The user's entry point (e.g., domains, IP addresses, and/or geolocation information such as country/city) from which a user logs in.
2. The login user and machine class.
3. Binaries, executables, processes, etc. a user launches.
4. Internal servers with which the user (or any of the user's processes, child processes, etc.) communicates, and external contacts (e.g., domains, IP addresses, and/or geolocation information such as country/city) with which the user communicates (i.e., transfers data).

In the event of a security breach, being able to concretely answer questions about such information can be very important. And, collectively, such information is useful in providing an end-to-end path (e.g., for performing investigations).

In the following example, suppose a user ("UserA") logs into a machine ("Machine01") from a first IP address ("IP01"). Machine01 is inside a datacenter. UserA then launches a script ("runnable.sh") on Machine01. From Machine01, UserA next logs into a second machine ("Machine02") via ssh, also as UserA, also within the datacenter. On Machine02, UserA again launches a script ("new_runnable.sh"). On Machine02, UserA then changes privilege, becoming root on Machine02. From Machine02, UserA (now as root) logs into a third machine ("Machine03") in the datacenter via ssh, as root on Machine03. As root on Machine03, the user executes a script ("collect_data.sh") on Machine03. The script internally communicates (as root) to a MySQL-based service internal to the datacenter, and downloads data from the MySQL-based service. Finally, as root on Machine03, the user externally communicates with a server outside the datacenter ("External01"), using a POST command. To summarize what has occurred, in this example, the source/entry point is IP01. Data is transferred to an external server External01. The machine performing the transfer to External01 is Machine03. The user transferring the data is "root" (on Machine03), while the actual user (hiding behind root) is UserA.

In the above scenario, the "original user" (ultimately responsible for transmitting data to External01) is UserA,

who logged in from IP01. Each of the processes ultimately started by UserA, whether started at the command line (tty) such as “runnable.sh” or started after an ssh connection such as “new_runnable.sh,” and whether as UserA, or as a subsequent identity, are all examples of child processes which can be arranged into a process hierarchy.

As previously mentioned, machines can be clustered together logically into machine clusters. One approach to clustering is to classify machines based on information such as the types of services they provide/binaries they have installed upon them/processes they execute. Machines sharing a given machine class (as they share common binaries/services/etc.) will behave similarly to one another. Each machine in a datacenter can be assigned to a machine cluster, and each machine cluster can be assigned an identifier (also referred to herein as a machine class). One or more tags can also be assigned to a given machine class (e.g., database_servers_west or prod_web_frontend). One approach to assigning a tag to a machine class is to apply term frequency analysis (e.g., TF/IDF) to the applications run by a given machine class, selecting as tags those most unique to the class. Data platform 12 can use behavioral baselines taken for a class of machines to identify deviations from the baseline (e.g., by a particular machine in the class).

FIG. 3S illustrates an example of a process for detecting anomalies. In various embodiments, process 392 is performed by data platform 12. As explained above, a given session will have an original user. And, each action taken by the original user can be tied back to the original user, despite privilege changes and/or lateral movement throughout a datacenter. Process 392 begins at 393 when log data associated with a user session (and thus an original user) is received. At 394, a logical graph is generated, using at least a portion of the collected data. When an anomaly is detected (395), it can be recorded, and as applicable, an alert is generated (396). The following are examples of graphs that can be generated (e.g., at 394), with corresponding examples of anomalies that can be detected (e.g., at 395) and alerted upon (e.g., at 396).

FIG. 4A illustrates a representation of an embodiment of an insider behavior graph. In the example of FIG. 4A, each node in the graph can be: (1) a cluster of users; (2) a cluster of launched processes; (3) a cluster of processes/servers running on a machine class; (4) a cluster of external IP addresses (of incoming clients); or (5) a cluster of external servers based on DNS/IP/etc. As depicted in FIG. 4A, graph data is vertically tiered into four tiers. Tier 0 (400) corresponds to entry point information (e.g., domains, IP addresses, and/or geolocation information) associated with a client entering the datacenter from an external entry point. Entry points are clustered together based on such information. Tier 1 (401) corresponds to a user on a machine class, with a given user on a given machine class represented as a node. Tier 2 (402) corresponds to launched processes, child processes, and/or interactive processes. Processes for a given user and having similar connectivity (e.g., sharing the processes they launch and the machines with which they communicate) are grouped into nodes. Finally, Tier 3 (403) corresponds to the services/servers/domains/IP addresses with which processes communicate. A relationship between the tiers can be stated as follows: Tier 0 nodes log in to tier 1 nodes. Tier 1 nodes launch tier 2 nodes. Tier 2 nodes connect to tier 3 nodes.

The inclusion of an original user in both Tier 1 and Tier 2 allows for horizontal tiering. Such horizontal tiering ensures that there is no overlap between any two users in Tier 1 and Tier 2. Such lack of overlap provides for faster

searching of an end-to-end path (e.g., one starting with a Tier 0 node and terminating at a Tier 3 node). Horizontal tiering also helps in establishing baseline insider behavior. For example, by building an hourly insider behavior graph, new edges/changes in edges between nodes in Tier 1 and Tier 2 can be identified. Any such changes correspond to a change associated with the original user. And, any such changes can be surfaced as anomalous and alerts can be generated.

As explained above, Tier 1 corresponds to a user (e.g., user “U”) logging into a machine having a particular machine class (e.g., machine class “M”). Tier 2 is a cluster of processes having command line similarity (e.g., CType “C”), having an original user “U,” and running as a particular effective user (e.g., user “U1”). The value of U1 may be the same as U (e.g., joe.smith in both cases), or the value of U1 may be different (e.g., U-joe.smith and U1=root). Thus, while an edge may be present from a Tier 1 node to a Tier 2 node, the effective user in the Tier 2 node may or may not match the original user (while the original user in the Tier 2 node will match the original user in the Tier 1 node).

A change from a user U into a user U1 can take place in a variety of ways. Examples include where U becomes U1 on the same machine (e.g., via su), and also where U sshes to other machine(s). In both situations, U can perform multiple changes, and can combine approaches. For example, U can become U1 on a first machine, ssh to a second machine (as U1), become U2 on the second machine, and ssh to a third machine (whether as user U2 or user U3). In various embodiments, the complexity of how user U ultimately becomes U3 (or U5, etc.) is hidden from a viewer of an insider behavior graph, and only an original user (e.g., U) and the effective user of a given node (e.g., U5) are depicted. As applicable (e.g., if desired by a viewer of the insider behavior graph), additional detail about the path (e.g., an end-to-end path of edges from user U to user U5) can be surfaced (e.g., via user interactions with nodes).

FIG. 4B illustrates an example of a portion of an insider behavior graph (e.g., as rendered in a web browser). In the example shown, node 405 (the external IP address, 52.32.40.231) is an example of a Tier 0 node, and represents an entry point into a datacenter. As indicated by directional arrows 406 and 407, two users, “aruneli_prod” and “harish_prod,” both made use of the source IP 52.32.40.231 when logging in between 5 pm and 6 pm on Sunday July 30 (408). Nodes 409 and 410 are examples of Tier 1 nodes, having aruneli_prod and harish_prod as associated respective original users. As previously mentioned, Tier 1 nodes correspond to a combination of a user and a machine class. In the example depicted in FIG. 4B, the machine class associated with nodes 409 and 410 is hidden from view to simplify visualization, but can be surfaced to a viewer of interface 404 (e.g., when the user clicks on node 409 or 410).

Nodes 414-423 are examples of Tier 2 nodes-processes that are launched by users in Tier 1 and their child, grand-child, etc. processes. Note that also depicted in FIG. 4B is a Tier 1 node 411 that corresponds to a user, “root,” that logged in to a machine cluster from within the datacenter (i.e., has an entry point within the datacenter). Nodes 425-1 and 425-2 are examples of Tier 3 nodes—internal/external IP addresses, servers, etc., with which Tier 2 nodes communicate.

In the example shown in FIG. 4B, a viewer of interface 404 has clicked on node 423. As indicated in region 426, the user running the marathon container is “root.” However, by following the directional arrows in the graph backwards from node 423 (i.e. from right to left), the viewer can

determine that the original user, responsible for node 423, is “aruneli_prod,” who logged into the datacenter from IP 52.32.40.231.

The following are examples of changes that can be tracked using an insider behavior graph model:

A user logs in from a new IP address.

A user logs in from a geolocation not previously used by that user.

A user logs into a new machine class.

A user launches a process not previously used by that user.

A user connects to an internal server to which the user has not previously connected.

An original user communicates with an external server (or external server known to be malicious) with which that user has not previously communicated.

A user communicates with an external server which has a geolocation not previously used by that user.

Such changes can be surfaced as alerts, e.g., to help an administrator determine when/what anomalous behavior occurs within a datacenter. Further, the behavior graph model can be used (e.g., during forensic analysis) to answer questions helpful during an investigation. Examples of such questions include:

Was there any new login activity (Tier 0) in the timeframe being investigated? As one example, has a user logged in from an IP address with unknown geolocation information? Similarly, has a user started communicating externally with a new Tier 3 node (e.g., one with unknown geolocation information).

Has there been any suspicious login activity (Tier 0) in the timeframe being investigated? As one example, has a user logged in from an IP address that corresponds to a known bad IP address as maintained by Threat aggregator 150? Similarly, has there been any suspicious Tier 3 activity?

Were any anomalous connections made within the datacenter during the timeframe being investigated? As one example, suppose a given user (“Frank”) typically enters a datacenter from a particular IP address (or range of IP addresses), and then connects to a first machine type (e.g., bastion), and then to a second machine type (e.g., database_prod). If Frank has directly connected to database_prod (instead of first going through bastion) during the timeframe, this can be surfaced using the insider graph.

Who is (the original user) responsible for running a particular process?

An example of an insider behavior graph being used in an investigation is depicted in FIGS. 4C and 4D. FIG. 4C depicts a baseline of behavior for a user, “Bill.” As shown in FIG. 4C, Bill typically logs into a datacenter from the IP address, 71.198.44.40 (427). He typically makes use of ssh (428), and sudo (429), makes use of a set of typical applications (430) and connects (as root) with the external service, api.lacework.net (431).

Suppose Bill’s credentials are compromised by a nefarious outsider (“Eve”). FIG. 4D depicts an embodiment of how the graph depicted in FIG. 4C would appear once Eve begins exfiltrating data from the datacenter. Eve logs into the datacenter (using Bill’s credentials) from 52.5.66.8 (432). As Bill, Eve escalates her privilege to root (e.g., via su), and then becomes a different user, Alex (e.g., via su alex). As Alex, Eve executes a script, “sneak.sh” (433), which launches another script, “post.sh” (434), which contacts external server 435 which has an IP address of 52.5.66.7, and transmits data to it. Edges 436-439 each represent changes in Bill’s behavior. As previously mentioned, such

changes can be detected as anomalies and associated alerts can be generated. As a first example, Bill logging in from an IP address he has not previously logged in from (436) can generate an alert. As a second example, while Bill does typically make use of sudo (429), he has not previously executed sneak.sh (433) or post.sh (434) and the execution of those scripts can generate alerts as well. As a third example, Bill has not previously communicated with server 435, and an alert can be generated when he does so (439). Considered individually, each of edges 436-439 may indicate nefarious behavior, or may be benign. As an example of a benign edge, suppose Bill begins working from a home office two days a week. The first time he logs in from his home office (i.e., from an IP address that is not 71.198.44.40), an alert can be generated that he has logged in from a new location. Over time, however, as Bill continues to log in from his home office but otherwise engages in typical activities, Bill’s graph will evolve to include logins from both 71.198.44.40 and his home office as baseline behavior. Similarly, if Bill begins using a new tool in his job, an alert can be generated the first time he executes the tool, but over time will become part of his baseline.

In some cases, a single edge can indicate a serious threat. For example, if server 432 (or 435) is included in a known bad IP listing, edge 436 (or 439) indicates compromise. An alert that includes an appropriate severity level (e.g., “threat level high”) can be generated. In other cases, a combination of edges could indicate a threat (where a single edge might otherwise result in a lesser warning). In the example shown in FIG. 4D, the presence of multiple new edges is indicative of a serious threat. Of note, even though “sneak.sh” and “post.sh” were executed by Alex, because data platform 12 also keeps track of an original user, the compromise of user B’s account will be discovered.

FIG. 4E illustrates a representation of an embodiment of a user login graph. In the example of FIG. 4E, tier 0 (440) clusters source IP addresses as belonging to a particular country (including an “unknown” country) or as a known bad IP. Tier 1 (441) clusters user logins, and tier 2 (442) clusters type of machine class into which a user is logging in. The user login graph tracks the typical login behavior of users. By interacting with a representation of the graph, answers to questions such as the following can be obtained:

Where is a user logging in from?

Have any users logged in from a known bad address?

Have any non-developer users accessed development machines?

Which machines does a particular user access?

Examples of alerts that can be generated using the user login graph include:

A user logs in from a known bad IP address.

A user logs in from a new country for the first time.

A new user logs into the datacenter for the first time.

A user accesses a machine class that the user has not previously accessed.

One way to track privilege changes in a datacenter is by monitoring a process hierarchy of processes. To help filter out noisy commands/processes such as “su -u,” the hierarchy of processes can be constrained to those associated with network activity. In a *nix system, each process has two identifiers assigned to it, a process identifier (PID) and a parent process identifier (PPID). When such a system starts, the initial process is assigned a PID 0. Each user process has a corresponding parent process.

Using techniques described herein, a graph can be constructed (also referred to herein as a privilege change graph) which models privilege changes. In particular, a graph can

be constructed which identifies where a process P1 launches a process P2, where P1 and P2 each have an associated user U1 and U2, with U1 being an original user, and U2 being an effective user. In the graph, each node is a cluster of processes (sharing a CType) executed by a particular (original) user. As all the processes in the cluster belong to the same user, a label that can be used for the cluster is the user's username. An edge in the graph, from a first node to a second node, indicates that a user of the first node changed its privilege to the user of the second node.

FIG. 4F illustrates an example of a privilege change graph. In the example shown in FIG. 4F, each node (e.g., nodes 444 and 445) represents a user. Privilege changes are indicated by edges, such as edge 446.

As with other graphs, anomalies in graph 443 can be used to generate alerts. Three examples of such alerts are as follows:

New user entering the datacenter. Any time a new user enters the datacenter and runs a process, the graph will show a new node, with a new CType. This indicates a new user has been detected within the datacenter. FIG. 4F is a representation of an example of an interface that depicts such an alert. Specifically, as indicated in region 447, an alert for the time period 1 pm-2 pm on June 8 was generated. The alert identifies that a new user, Bill (448) executed a process.

Privilege change. As explained above, a new edge, from a first node (user A) to a second node (user B) indicates that user A has changed privilege to user B.

Privilege escalation. Privilege escalation is a particular case of privilege change, in which the first user becomes root.

An example of an anomalous privilege change and an example of an anomalous privilege escalation are each depicted in graph 450 of FIG. 4G. In particular, as indicated in region 451, two alerts for the time period 2 pm-3 pm on June 8 were generated (corresponding to the detection of the two anomalous events). In region 452, root has changed privilege to the user "daemon," which root has not previously done. This anomaly is indicated to the user by highlighting the daemon node (e.g., outlining it in a particular color, e.g., red). As indicated by edge 453, Bill has escalated his privilege to the user root (which can similarly be highlighted in region 454). This action by Bill represents a privilege escalation.

An Extensible query interface for dynamic data compositions and filter applications will now be described.

As described herein, datacenters are highly dynamic environments. And, different customers of data platform 12 (e.g., entity A vs. entity B) may have different/disparate needs/requirements of data platform 12, e.g., due to having different types of assets, different applications, etc. Further, as time progresses, new software tools will be developed, new types of anomalous behavior will be possible (and should be detectable), etc. In various embodiments, data platform 12 makes use of predefined relational schema (including by having different predefined relational schema for different customers). However, the complexity and cost of maintaining/updating such predefined relational schema can rapidly become problematic-particularly where the schema includes a mix of relational, nested, and hierarchical (graph) datasets. In other embodiments, the data models and filtering applications used by data platform 12 are extensible. As will be described in more detail below, in various embodiments, data platform 12 supports dynamic query generation by automatic discovery of join relations via static or dynamic filtering key specifications among composable data sets.

This allows a user of data platform 12 to be agnostic to modifications made to existing data sets as well as creation of new data sets. The extensible query interface also provides a declarative and configurable specification for optimizing internal data generation and derivations.

As will also be described in more detail below, data platform 12 is configured to dynamically translate user interactions (e.g., received via web app 120) into SQL queries (and without the user needing to know how to write queries). Such queries can then be performed (e.g., by query service 166) against any compatible backend (e.g., data store 30).

FIG. 4H illustrates an example of a user interacting with a portion of an interface. When a user visits data platform 12 (e.g., via web app 120 using a browser), data is extracted from data store 30 as needed (e.g., by query service 166), to provide the user with information, such as the visualizations depicted variously herein. As the user continues to interact with such visualizations (e.g., clicking on graph nodes, entering text into search boxes, navigating between tabs (e.g., tab 455 vs. 465)), such interactions act as triggers that cause query service 166 to continue to obtain information from data store 30 as needed (and as described in more detail below).

In the example shown in FIG. 4H, user A is viewing a dashboard that provides various information about entity A users (455), during the time period March 2 at midnight-March 25 at 7 pm (which she selected by interacting with region 456). Various statistical information is presented to user A in region 457. Region 458 presents a timeline of events that occurred during the selected time period. User A has opted to list only the critical, high, and medium events during the time period by clicking on the associated boxes (459-461). A total of 55 low severity, and 155 info-only events also occurred during the time period. Each time user A interacts with an element in FIG. 4H (e.g., clicks on box 461, clicks on link 464-1, or clicks on tab 465), her actions are translated/formalized into filters on the data set and used to dynamically generate SQL queries. The SQL queries are generated transparently to user A (and also to a designer of the user interface shown in FIG. 4H).

User A notes in the timeline (462) that a user, Harish, connected to a known bad server (examplebad.com) using wget, an event that has a critical severity level. User A can click on region 463 to expand details about the event inline (which will display, for example, the text "External connection made to known bad host examplebad.com at port 80 from application 'wget' running on host dev1.lacework.internal as user harish") directly below timeline 462. User A can also click on link 464-1, which will take her to a dossier for the event (depicted in FIG. 4I). As will be described in more detail below, a dossier is a template for a collection of visualizations.

As shown in interface 466, the event of Harish using wget to contact examplebad.com on March 16 was assigned an event ID of 9291 by data platform 12 (467). For convenience to user A, the event is also added to her dashboard in region 476 as a bookmark (468). A summary of the event is depicted in region 469. By interacting with boxes shown in region 470, user A can see a timeline of related events. In this case, user A has indicated that she would like to see other events involving the wget application (by clicking box 471). Events of critical and medium security involving wget occurred during the one hour window selected in region 472.

Region 473 automatically provides user A with answers to questions that may be helpful to have answers to while investigating event 9291. If user A clicks on any of the links

in the event description (474), she will be taken to a corresponding dossier for the link. As one example, suppose user A clicks on link 475. She will then be presented with interface 477 shown in FIG. 4J.

Interface 477 is an embodiment of a dossier for a domain. In this example, the domain is “examplebad.com,” as shown in region 478. Suppose user A would like to track down more information about interactions entity A resources have made with examplebad.com between January 1 and March 20. She selects the appropriate time period in region 479 and information in the other portions of interface 477 automatically update to provide various information corresponding to the selected time frame. As one example, user A can see that contact was made with examplebad.com a total of 17 times during the time period (480), as well as a list of each contact (481). Various statistical information is also included in the dossier for the time period (482). If she scrolls down in interface 477, user A will be able to view various polygraphs associated with examplebad.com, such as an application-communication polygraph (483).

Data stored in data store 30 can be internally organized as an activity graph. In the activity graph, nodes are also referred to as Entities. Activities generated by Entities are modeled as directional edges between nodes. Thus, each edge is an activity between two Entities. One example of an Activity is a “login” Activity, in which a user Entity logs into a machine Entity (with a directed edge from the user to the machine). A second example of an Activity is a “launch” Activity, in which a parent process launches a child process (with a directed edge from the parent to the child). A third example of an Activity is a “DNS query” Activity, in which either a process or a machine performs a query (with a directed edge from the requestor to the answer, e.g., an edge from a process to www.example.com). A fourth example of an Activity is a network “connected to” Activity, in which processes, IP addresses, and listen ports can connect to each other (with a directed edge from the initiator to the server).

As will be described in more detail below, query service 166 provides either relational views or graph views on top of data stored in data store 30. Typically, a user will want to see data filtered using the activity graph. For example, if an entity was not involved in an activity in a given time period, that entity should be filtered out of query results. Thus, a request to show “all machines” in a given time frame will be interpreted as “show distinct machines that were active” during the time frame.

Query service 166 relies on three main data model elements: fields, entities, and filters. As used herein, a field is a collection of values with the same type (logical and physical). A field can be represented in a variety of ways, including: 1. a column of relations (table/view), 2. a return field from another entity, 3. an SQL aggregation (e.g., COUNT, SUM, etc.), 4. an SQL expression with the references of other fields specified, and 5. a nested field of a JSON object. As viewed by query service 166, an entity is a collection of fields that describe a data set. The data set can be composed in a variety of ways, including: 1. a relational table, 2. a parameterized SQL statement, 3. DynamicSQL created by a Java function, and 4. join/project/aggregate/subclass of other entities. Some fields are common for all entities. One example of such a field is a “first observed” timestamp (when first use of the entity was detected). A second example of such a field is the entity classification type (e.g., one of: 1. Machine (on which an agent is installed), 2. Process, 3. Binary, 4. UID, 5. IP, 6. DNS Information, 7. ListenPort, and 8. PType). A third example of such a field is a “last observed” timestamp.

A filter is an operator that: 1. takes an entity and field values as inputs, 2. a valid SQL expression with specific reference(s) of entity fields, or 3. is a conjunct/disjunct of filters. As will be described in more detail below, filters can be used to filter data in various ways, and limit data returned by query service 166 without changing the associated data set.

As mentioned above, a dossier is a template for a collection of visualizations. Each visualization (e.g., the box including chart 484) has a corresponding card, which identifies particular target information needed (e.g., from data store 30) to generate the visualization. In various embodiments, data platform 12 maintains a global set of dossiers/cards. Users of data platform 12 such as user A can build their own dashboard interfaces using preexisting dossiers/cards as components, and/or they can make use of a default dashboard (which incorporates various of such dossiers/cards).

A JSON file can be used to store multiple cards (e.g., as part of a query service catalog). A particular card is represented by a single JSON object with a unique name as a field name. Each card may be described by the following named fields:

TYPE: the type of the card. Example values include:
Entity (the default type)

SQL
Filters
DynamicSQL
graphFilter
graph
Function
Template

PARAMETERS: a JSON array object that contains an array of parameter objects with the following fields:

name (the name of the parameter)
required (a Boolean flag indicating whether the parameter is required or not)
default (a default value of the parameter)
props (a generic JSON object for properties of the parameter. Possible values are:
“utype” (a user defined type), and “scope” (an optional property to configure a namespace of the parameter))
value (a value for the parameter-non-null to override the default value defined in nested source entities)

SOURCES: a JSON array object explicitly specifying references of input entities. Each source reference has the following attributes:

name (the card/entity name or fully-qualified Table name)
type (required for base Table entity)
alias (an alias to access this source entity in other fields (e.g., returns, filters, groups, etc))

RETURNS: a required JSON array object of a return field object. A return field object can be described by the following attributes:

field (a valid field name from a source entity)
expr (a valid SQL scalar expression. References to input fields of source entities are specified in the format of #{Entity.Field}. Parameters can also be used in the expression in the format of \${ParameterName})
type (the type of field, which is required for return fields specified by expr. It is also required for all return fields of an Entity with an SQL type)
alias (the unique alias for return field)
aggr (possible aggregations are: COUNT, COUNT_DISTINCT, DISTINCT, MAX, MIN, AVG, SUM, FIRST_VALUE, LAST_VALUE)

case (JSON array object represents conditional expressions “when” and “expr”)
 fieldsFrom, and, except (specification for projections from a source entity with excluded fields)
 props (general JSON object for properties of the return field. Possible properties include: “filterGroup,” “title,” “format,” and “utype”)
 PROPS: generic JSON objects for other entity properties
 SQL: a JSON array of string literals for SQL statements. Each string literal can contain parameterized expressions $\{ParameterName\}$ and/or composable entity by $\{EntityName\}$ GRAPH: required for graph entity. Has the following required fields:

source (including “type,” “props,” and “keys”)
 target (including “type,” “props,” and “keys”)
 edge (including “type” and “props”)
 JOINS: a JSON array of join operators. Possible fields for a join operator include:

type (possible join types include: “loj”-Left Outer Join, “join”-Inner Join, “in”-Semi Join, “implicit”-Implicit Join)
 left (a left hand side field of join)
 right (a right hand side field of join)
 keys (key columns for multi-way joins)
 order (a join order of multi-way joins)
 FKEYS: a JSON array of FilterKey(s). The fields for a FilterKey are:

type (type of FilterKey)
 fieldRefs (reference(s) to return fields of an entity defined in the sources field)
 alias (an alias of the FilterKey, used in implicit join specification)
 FILTERS: a JSON array of filters (conjunct). Possible fields for a filter include:

type (types of filters, including: “eq”—equivalent to SQL=, “ne”—equivalent to SQL<>, “ge”—equivalent to SQL>=, “gt”—equivalent to SQL>, “le”—equivalent to SQL<=, “lt”—equivalent to SQL<, “like”—equivalent to SQL LIKE, “not_like”—equivalent to SQL NOT LIKE, “rlike”—equivalent to SQL RLIKE (Snowflake specific), “not_rlike”—equivalent to SQL NOT RLIKE (Snowflake specific), “in”—equivalent to SQL IN, “not_in”—equivalent to SQL NOT IN)
 expr (generic SQL expression)
 field (field name)
 value (single value)
 values (for both IN and NOT IN)

ORDERS: a JSON array of ORDER BY for returning fields. Possible attributes for the ORDER BY clause include:
 field (field ordinal index (1 based) or field alias)
 order (asc/desc, default is ascending order)
 GROUPS: a JSON array of GROUP BY for returning fields. Field attributes are:

field (ordinal index (1 based) or alias from the return fields)
 LIMIT: a limit for the number of records to be returned
 OFFSET: an offset of starting position of returned data. Used in combination with limit for pagination.

Suppose customers of data platform 12 (e.g., entity A and entity B) request new data transformations or a new aggregation of data from an existing data set (as well as a corresponding visualization for the newly defined data set). As mentioned above, the data models and filtering applications used by data platform 12 are extensible. Thus, two example scenarios of extensibility are (1) extending the filter data set, and (2) extending a FilterKey in the filter data set.

Data platform 12 includes a query service catalog that enumerates cards available to users of data platform 12. New cards can be included for use in data platform 12 by being added to the query service catalog (e.g., by an operator of data platform 12). For reusability and maintainability, a single external-facing card (e.g., available for use in a dossier) can be composed of multiple (nested) internal cards. Each newly added card (whether external or internal) will also have associated FilterKey(s) defined. A user interface (UI) developer can then develop a visualization for the new data set in one or more dossier templates. The same external card can be used in multiple dossier templates, and a given external card can be used multiple times in the same dossier (e.g., after customization). Examples of external card customization include customization via parameters, ordering, and/or various mappings of external data fields (columns).

As mentioned above, a second extensibility scenario is one in which a FilterKey in the filter data set is extended (i.e., existing template functions are used to define a new data set). As also mentioned above, data sets used by data platform 12 are composable/reusable/extensible, irrespective of whether the data sets are relational or graph data sets. One example data set is the User Tracking polygraph, which is generated as a graph data set (comprising nodes and edges). Like other polygraphs, User Tracking is an external data set that can be visualized both as a graph (via the nodes and edges) and can also be used as a filter data set for other cards, via the cluster identifier (CID) field.

As mentioned above, as users such as user A navigate through/interact with interfaces provided by data platform 12 (e.g., as shown in FIG. 4H), such interactions trigger query service 166 to generate and perform queries against data store 30. Dynamic composition of filter datasets can be implemented using FilterKeys and FilterKey Types. A FilterKey can be defined as a list of columns and/or fields in a nested structure (e.g., JSON). Instances of the same FilterKey Type can be formed as an Implicit Join Group. The same instance of a FilterKey can participate in different Implicit Join Groups. A list of relationships among all possible Implicit Join Groups is represented as a Join graph for the entire search space to create a final data filter set by traversing edges and producing Join Path(s).

Each card (e.g., as stored in the query service catalog and used in a dossier) can be introspected by a/card/describe/CardID REST request.

At runtime (e.g., whenever it receives a request from web app 120), query service 166 parses the list of implicit joins and creates a Join graph to manifest relationships of FilterKeys among Entities. A Join graph (an example of which is depicted in FIG. 4K) comprises a list of Join Link(s). A Join Link represents each implicit join group by the same FilterKey type. A Join Link maintains a reverse map (Entity-to-FilterKey) of FilterKeys and their Entities. As previously mentioned, Entities can have more than one FilterKey defined. The reverse map guarantees one FilterKey per Entity can be used for each JoinLink. Each JoinLink also maintains a list of entities for the priority order of joins. Each JoinLink is also responsible for creating and adding directional edge(s) to graphs. An edge represents a possible join between two Entities.

At runtime, each Implicit Join uses the Join graph to find all possible join paths. The search of possible join paths starts with the outer FilterKey of an implicit join. One approach is to use a shortest path approach, with breadth first traversal and subject to the following criteria:

Use the priority order list of Join Links for all entities in the same implicit join group.

Stop when a node (Entity) is reached which has local filter(s).

Include all join paths at the same level (depth).

Exclude join paths based on the predefined rules (path of edges).

FIG. 4L illustrates an example of a process for dynamically generating and executing a query. In various embodiments, process 485 is performed by data platform 12. The process begins at 486 when a request is received to filter information associated with activities within a network environment. One example of such a request occurs in response to user A clicking on tab 465. Another example of such a request occurs in response to user A clicking on link 464-1. Yet another example of such a request occurs in response to user A clicking on link 464-2 and selecting (e.g., from a dropdown) an option to filter (e.g., include, exclude) based on specific criteria that she provides (e.g., an IP address, a username, a range of criteria, etc.).

At 487, a query is generated based on an implicit join. One example of processing that can be performed at 487 is as follows. As explained above, one way dynamic composition of filter datasets can be implemented is by using FilterKeys and FilterKey Types. And, instances of the same FilterKey Type can be formed as an Implicit Join Group. A Join graph for the entire search space can be constructed from a list of all relationships among all possible Join Groups. And, a final data filter set can be created by traversing edges and producing one or more Join Paths. Finally, the shortest path in the join paths is used to generate an SQL query string.

One approach to generating an SQL query string is to use a query building library (authored in an appropriate language such as Java). For example, a common interface “sqlGen” may be used in conjunction with process 485 is as follows. First, a card/entity is composed by a list of input cards/entities, where each input card recursively is composed by its own list of input cards. This nested structure can be visualized as a tree of query blocks (SELECT) in standard SQL constructs. SQL generation can be performed as the traversal of the tree from root to leaf entities (top-down), calling the sqlGen of each entity. Each entity can be treated as a subclass of the Java class (Entity). An implicit join filter (EntityFilter) is implemented as a subclass of Entity, similar to the right hand side of a SQL semi-join operator. Unlike the static SQL semi-join construct, it is conditionally and recursively generated even if it is specified in the input sources of the JSON specification. Another recursive interface can also be used in conjunction with process 485, preSQLGen, which is primarily the entry point for EntityFilter to run a search and generate nested implicit join filters. During preSQLGen recursive invocations, the applicability of implicit join filters is examined and pushed down to its input subquery list. Another top-down traversal, pullUp-Cachable, can be used to pull up common sub-query blocks, including those dynamically generated by preSQLGen, such that SELECT statements of those cacheable blocks are generated only once at top-level WITH clauses. A recursive interface, sqlWith, is used to generate nested subqueries inside WITH clauses. The recursive calls of a sqlWith function can generate nested WITH clauses as well. An sqlFrom function can be used to generate SQL FROM clauses by referencing those subquery blocks in the WITH clauses. It also produces INNER/OUTER join operators based on the joins in the specification. Another recursive interface, sqlWhere, can be used to generate conjuncts and disjuncts of local predicates and semi-join predicates based on implicit join transformations. Further, sqlProject, sql-

GroupBy, sqlOrderBy, and sqlLimitOffset can respectively be used to translate the corresponding directives in JSON spec to SQL SELECT list, GROUP BY, ORDER BY, and LIMIT/OFFSET clauses.

Returning to process 485, at 488, the query (generated at 487) is used to respond to the request. As one example of the processing performed at 488, the generated query is used to query data store 30 and provide (e.g., to web app 120) fact data formatted in accordance with a schema (e.g., as associated with a card associated with the request received at 486).

Although the examples described herein largely relate to embodiments where data is collected from agents and ultimately stored in a data store such as those provided by Snowflake, in other embodiments data that is collected from agents and other sources may be stored in different ways. For example, data that is collected from agents and other sources may be stored in a data warehouse, data lake, data mart, and/or any other data store.

A data warehouse may be embodied as an analytic database (e.g., a relational database) that is created from two or more data sources. Such a data warehouse may be leveraged to store historical data, often on the scale of petabytes. Data warehouses may have compute and memory resources for running complicated queries and generating reports. Data warehouses may be the data sources for business intelligence (“BI”) systems, machine learning applications, and/or other applications. By leveraging a data warehouse, data that has been copied into the data warehouse may be indexed for good analytic query performance, without affecting the write performance of a database (e.g., an Online Transaction Processing (“OLTP”) database). Data warehouses also enable the joining data from multiple sources for analysis. For example, a sales OLTP application probably has no need to know about the weather at various sales locations, but sales predictions could take advantage of that data. By adding historical weather data to a data warehouse, it would be possible to factor it into models of historical sales data.

Data lakes, which store files of data in their native format, may be considered as “schema on read” resources. As such, any application that reads data from the lake may impose its own types and relationships on the data. Data warehouses, on the other hand, are “schema on write,” meaning that data types, indexes, and relationships are imposed on the data as it is stored in the EDW. “Schema on read” resources may be beneficial for data that may be used in several contexts and poses little risk of losing data. “Schema on write” resources may be beneficial for data that has a specific purpose, and good for data that must relate properly to data from other sources. Such data stores may include data that is encrypted using homomorphic encryption, data encrypted using privacy-preserving encryption, smart contracts, non-fungible tokens, decentralized finance, and other techniques.

Data marts may contain data oriented towards a specific business line whereas data warehouses contain enterprise-wide data. Data marts may be dependent on a data warehouse, independent of the data warehouse (e.g., drawn from an operational database or external source), or a hybrid of the two. In embodiments described herein, different types of data stores (including combinations thereof) may be leveraged. Such data stores may be proprietary or may be embodied as vendor provided products or services such as, for example, Google BigQuery, Druid, Amazon Redshift, IBM Db2, Dremio, Databricks Lakehouse Platform, Cloudera, Azure Synapse Analytics, and others.

The deployments (e.g., a customer’s cloud deployment) that are analyzed, monitored, evaluated, or otherwise

observed by the systems described herein (e.g., systems that include components such as the platform 12 of FIG. 1D, the data collection agents described herein, and/or other components) may be provisioned, deployed, and/or managed using infrastructure as code ('IaC'). IaC involves the managing and/or provisioning of infrastructure through code instead of through manual processes. With IaC, configuration files may be created that include infrastructure specifications. IaC can be beneficial as configurations may be edited and distributed, while also ensuring that environments are provisioned in a consistent manner. IaC approaches may be enabled in a variety of ways including, for example, using IaC software tools such as Terraform by HashiCorp. Through the usage of such tools, users may define and provide data center infrastructure using JavaScript Object Notation ('JSON'), YAML, proprietary formats, or some other format. In some embodiments, the configuration files may be used to emulate a cloud deployment for the purposes of analyzing the emulated cloud deployment using the systems described herein. Likewise, the configuration files themselves may be used as inputs to the systems described herein, such that the configuration files may be inspected to identify vulnerabilities, misconfigurations, violations of regulatory requirements, or other issues. In fact, configuration files for multiple cloud deployments may even be used by the systems described herein to identify best practices, to identify configuration files that deviate from typical configuration files, to identify configuration files with similarities to deployments that have been determined to be deficient in some way, or the configuration files may be leveraged in some other ways to detect vulnerabilities, misconfigurations, violations of regulatory requirements, or other issues prior to deploying an infrastructure that is described in the configuration files. In some embodiments the techniques described herein may be used in multi-cloud, multi-tenant, cross-cloud, cross-tenant, cross-user, industry cloud, digital platform, and other scenarios depending on specific need or situation.

In some embodiments, the deployments that are analyzed, monitored, evaluated, or otherwise observed by the systems described herein (e.g., systems that include components such as the platform 12 of FIG. 1D, the data collection agents described herein, and/or other components) may be monitored to determine the extent to which a particular component has experienced "drift" relative to its associated IaC configuration. Discrepancies between how cloud resources were defined in an IaC configuration file and how they are currently configured in runtime may be identified and remediation workflows may be initiated to generate an alert, reconfigure the deployment, or take some other action. Such discrepancies may occur for a variety of reasons. Such discrepancies may occur, for example, due to maintenance operations being performed, due to incident response tasks being carried out, or for some other reason. Readers will appreciate that while IaC helps avoid initial misconfigurations of a deployment by codifying and enforcing resource creation, resource configuration, security policies, and so on, the systems described herein may prevent unwanted drift from occurring during runtime and after a deployment has been created in accordance with an IaC configuration.

In some embodiments, the deployments (e.g., a customer's cloud deployment) that are analyzed, monitored, evaluated, or otherwise observed by the systems described herein (e.g., systems that include components such as the platform 12 of FIG. 1D, the data collection agents described herein, and/or other components) may also be provisioned, deployed, and/or managed using security as code ('SaC').

SaC extends IaC concepts by defining cybersecurity policies and/or standards programmatically, so that the policies and/or standards can be referenced automatically in the configuration scripts used to provision cloud deployments. Stated differently, SaC can automate policy implementation and cloud deployments may even be compared with the policies to prevent "drift." For example, if a policy is created where all personally identifiable information ('PII') or personal health information ('PHI') must be encrypted when it is stored, that policy is translated into a process that is automatically launched whenever a developer submits code, and code that violates the policy may be automatically rejected.

In some embodiments, SaC may be implemented by initially classifying workloads (e.g., by sensitivity, by criticality, by deployment model, by segment). Policies that can be instantiated as code may subsequently be designed. For example, compute-related policies may be designed, access-related policies may be designed, application-related policies may be designed, network-related policies may be designed, data-related policies may be designed, and so on. Security as code may then be instantiated through architecture and automation, as successful implementation of SaC can benefit from making key architectural-design decisions and executing the right automation capabilities. Next, operating model protections may be built and supported. For example, an operating model may "shift left" to maximize self-service and achieve full-life-cycle security automation (e.g., by standardizing common development toolchains, CI/CD pipelines, and the like). In such an example, security policies and access controls may be part of the pipeline, automatic code review and bug/defect detection may be performed, automated build processes may be performed, vulnerability scanning may be performed, checks against a risk-control framework may be made, and other tasks may be performed all before deploying an infrastructure or components thereof.

The systems described herein may be useful in analyzing, monitoring, evaluating, or otherwise observing a GitOps environment. In a GitOps environment, Git may be viewed as the one and only source of truth. As such, GitOps may require that the desired state of infrastructure (e.g., a customer's cloud deployment) be stored in version control such that the entire audit trail of changes to such infrastructure can be viewed or audited. In a GitOps environment, all changes to infrastructure are embodied as fully traceable commits that are associated with committer information, commit IDs, time stamps, and/or other information. In such an embodiment, both an application and the infrastructure (e.g., a customer's cloud deployment) that supports the execution of the application are therefore versioned artifacts and can be audited using the gold standards of software development and delivery. Readers will appreciate that while the systems described herein are described as analyzing, monitoring, evaluating, or otherwise observing a GitOps environment, in other embodiments other source control mechanisms may be utilized for creating infrastructure, making changes to infrastructure, and so on. In these embodiments, the systems described herein may similarly be used for analyzing, monitoring, evaluating, or otherwise observing such environments.

As described in other portions of the present disclosure, the systems described herein may be used to analyze, monitor, evaluate, or otherwise observe a customer's cloud deployment. While securing traditional datacenters requires managing and securing an IP-based perimeter with networks and firewalls, hardware security modules ('HSMs'), security information and event management ('SIEM') technologies,

and other physical access restrictions, such solutions are not particularly useful when applied to cloud deployments. As such, the systems described herein may be configured to interact with and even monitor other solutions that are appropriate for cloud deployments such as, for example, “zero trust” solutions.

A zero trust security model (a.k.a., zero trust architecture) describes an approach to the design and implementation of IT systems. A primary concept behind zero trust is that devices should not be trusted by default, even if they are connected to a managed corporate network such as the corporate LAN and even if they were previously verified. Zero trust security models help prevent successful breaches by eliminating the concept of trust from an organization’s network architecture. Zero trust security models can include multiple forms of authentication and authorization (e.g., machine authentication and authorization, human/user authentication and authorization) and can also be used to control multiple types of accesses or interactions (e.g., machine-to-machine access, human-to-machine access).

In some embodiments, the systems described herein may be configured to interact with zero trust solutions in a variety of ways. For example, agents that collect input data for the systems described herein (or other components of such systems) may be configured to access various machines, applications, data sources, or other entity through a zero trust solution, especially where local instances of the systems described herein are deployed at edge locations. Likewise, given that zero trust solutions may be part of a customer’s cloud deployment, the zero trust solution itself may be monitored to identify vulnerabilities, anomalies, and so on. For example, network traffic to and from the zero trust solution may be analyzed, the zero trust solution may be monitored to detect unusual interactions, log files generated by the zero trust solution may be gathered and analyzed, and so on.

In some embodiments, the systems described herein may leverage various tools and mechanisms in the process of performing its primary tasks (e.g., monitoring a cloud deployment). For example, Linux eBPF is mechanism for writing code to be executed in the Linux kernel space. Through the usage of eBPF, user mode processes can hook into specific trace points in the kernel and access data structures and other information. For example, eBPF may be used to gather information that enables the systems described herein to attribute the utilization of networking resources or network traffic to specific processes. This may be useful in analyzing the behavior of a particular process, which may be important for observability/SIEM.

The systems described may be configured to collect security event logs (or any other type of log or similar record of activity) and telemetry in real time for threat detection, for analyzing compliance requirements, or for other purposes. In such embodiments, the systems described herein may analyze telemetry in real time (or near real time), as well as historical telemetry, to detect attacks or other activities of interest. The attacks or activities of interest may be analyzed to determine their potential severity and impact on an organization. In fact, the attacks or activities of interest may be reported, and relevant events, logs, or other information may be stored for subsequent examination.

In one embodiment, systems described herein may be configured to collect security event logs (or any other type of log or similar record of activity) and telemetry in real time to provide customers with a SIEM or SIEM-like solution. SIEM technology aggregates event data produced by security devices, network infrastructure, systems, applications,

or other source. Centralizing all of the data that may be generated by a cloud deployment may be challenging for a traditional SIEM, however, as each component in a cloud deployment may generate log data or other forms of machine data, such that the collective amount of data that can be used to monitor the cloud deployment can grow to be quite large. A traditional SIEM architecture, where data is centralized and aggregated, can quickly result in large amounts of data that may be expensive to store, process, retain, and so on. As such, SIEM technologies may frequently be implemented such that silos are created to separate the data.

In some embodiments of the present disclosure, data that is ingested by the systems described herein may be stored in a cloud-based data warehouse such as those provided by Snowflake and others. Given that companies like Snowflake offer data analytics and other services to operate on data that is stored in their data warehouses, in some embodiments one or more of the components of the systems described herein may be deployed in or near Snowflake as part of a secure data lake architecture (a.k.a., a security data lake architecture, a security data lake/warehouse). In such an embodiment, components of the systems described herein may be deployed in or near Snowflake to collect data, transform data, analyze data for the purposes of detecting threats or vulnerabilities, initiate remediation workflows, generate alerts, or perform any of the other functions that can be performed by the systems described herein. In such embodiments, data may be received from a variety of sources (e.g., EDR or EDR-like tools that handle endpoint data, cloud access security broker (‘CASB’) or CASB-like tools that handle data describing interactions with cloud applications, Identity and Access Management (‘IAM’) or IAM-like tools, and many others), normalized for storage in a data warehouse, and such normalized data may be used by the systems described herein. In fact, the systems described herein may actually implement the data sources (e.g., an EDR tool, a CASB tool, an IAM tool) described above.

In some embodiments one data source that is ingested by the systems described herein is log data, although other forms of data such as network telemetry data (flows and packets) and/or many other forms of data may also be utilized. In some embodiments, event data can be combined with contextual information about users, assets, threats, vulnerabilities, and so on, for the purposes of scoring, prioritization and expediting investigations. In some embodiments, input data may be normalized, so that events, data, contextual information, or other information from disparate sources can be analyzed more efficiently for specific purposes (e.g., network security event monitoring, user activity monitoring, compliance reporting). The embodiments described here offer real-time analysis of events for security monitoring, advanced analysis of user and entity behaviors, querying and long-range analytics for historical analysis, other support for incident investigation and management, reporting (for compliance requirements, for example), and other functionality.

In some embodiments, the systems described herein may be part of an application performance monitoring (‘APM’) solution. APM software and tools enable the observation of application behavior, observation of its infrastructure dependencies, observation of users and business key performance indicators (‘KPIs’) throughout the application’s life cycle, and more. The applications being observed may be developed internally, as packaged applications, as software as a service (‘SaaS’), or embodied in some other ways. In such

embodiments, the systems described herein may provide one or more of the following capabilities:

The ability to operate as an analytics platform that ingests, analyzes, and builds context from traces, metrics, logs, and other sources.

Automated discovery and mapping of an application and its infrastructure components.

Observation of an application's complete transactional behavior, including interactions over a data communications network.

Monitoring of applications running on mobile (native and browser) and desktop devices.

Identification of probable root causes of an application's performance problems and their impact on business outcomes.

Integration capabilities with automation and service management tools.

Analysis of business KPIs and user journeys (for example, login to check-out).

Domain-agnostic analytics capabilities for integrating data from third-party sources.

Endpoint monitoring to understand the user experience and its impact on business outcomes.

Support for virtual desktop infrastructure ('VDI') monitoring.

In embodiments where the systems described herein are used for APM, some components of the system may be modified, other components may be added, some components may be removed, and other components may remain the same. In such an example, similar mechanisms as described elsewhere in this disclosure may be used to collect information from the applications, network resources used by the application, and so on. The graph based modelling techniques may also be leveraged to perform some of the functions mentioned above, or other functions as needed.

In some embodiments, the systems described herein may be part of a solution for developing and/or managing artificial intelligence ('AI') or machine learning ('ML') applications. For example, the systems described herein may be part of an AutoML tool that automate the tasks associated with developing and deploying ML models. In such an example, the systems described herein may perform various functions as part of an AutoML tool such as, for example, monitoring the performance of a series of processes, micro-services, and so on that are used to collectively form the AutoML tool. In other embodiments, the systems described herein may perform other functions as part of an AutoML tool or may be used to monitor, analyze, or otherwise observe an environment that the AutoML tool is deployed within.

In some embodiments, the systems described herein may be used to manage, analyze, or otherwise observe deployments that include other forms of AI/ML tools. For example, the systems described herein may manage, analyze, or otherwise observe deployments that include AI services. AI services are, like other resources in an as-a-service model, ready-made models and AI applications that are consumable as services and made available through APIs. In such an example, rather than using their own data to build and train models for common activities, organizations may access pre-trained models that accomplish specific tasks. Whether an organization needs natural language processing ('NLP'), automatic speech recognition ('ASR'), image recognition, or some other capability, AI services simply plug-and-play into an application through an API. Likewise, the systems described herein may be used to manage, analyze, or otherwise observe deployments that include other forms of

AI/ML tools such as Amazon Sagemaker (or other cloud machine-learning platform that enables developers to create, train, and deploy ML models) and related services such as Data Wrangler (a service to accelerate data prep for ML) and Pipelines (a CI/CD service for ML).

In some embodiments, the systems described herein may be used to manage, analyze, or otherwise observe deployments that include various data services. For example, data services may include secure data sharing services, data marketplace services, private data exchanges services, and others. Secure data sharing services can allow access to live data from its original location, where those who are granted access to the data simply reference the data in a controlled and secure manner, without latency or contention from concurrent users. Because changes to data are made to a single version, data remains up-to-date for all consumers, which ensures data models are always using the latest version of such data. Data marketplace services operate as a single location to access live, ready-to-query data (or data that is otherwise ready for some other use). A data marketplace can even include a "feature stores," which can allow data scientists to repurpose existing work. For example, once a data scientist has converted raw data into a metric (e.g., costs of goods sold), this universal metric can be found quickly and used by other data scientists for quick analysis against that data.

In some embodiments, the systems described herein may be used to manage, analyze, or otherwise observe deployments that include distributed training engines or similar mechanisms such as, for example, such as tools built on Dask. Dask is an open source library for parallel computing that is written in Python. Dask is designed to enable data scientists to improve model accuracy faster, as Dask enables data scientists can do everything in Python end-to-end, which means that they no longer need to convert their code to execute in environments like Apache Spark. The result is reduced complexity and increased efficiency. The systems described herein may also be used to manage, analyze, or otherwise observe deployments that include technologies such as RAPIDS (an open source Python framework which is built on top of Dask). RAPIDS optimizes compute time and speed by providing data pipelines and executing data science code entirely on graphics processing units (GPUs) rather than CPUs. Multi-cluster, shared data architecture, DataFrames, Java user-defined functions (UDF) are supported to enable trained models to run within a data warehouse.

In some embodiments, the systems described herein may be leveraged for the specific use case of detecting and/or remediating ransomware attacks and/or other malicious action taken with respect to data, systems, and/or other resources associated with one or more entities.

Ransomware is a type of malware from cryptovirology that threatens to publish the victim's data or perpetually block access to such data unless a ransom is paid. In such embodiments, ransomware attacks may be carried out in a manner such that patterns (e.g., specific process-to-process communications, specific data access patterns, unusual amounts of encryption/re-encryption activities) emerge, where the systems described herein may monitor for such patterns. Alternatively, ransomware attacks may involve behavior that deviates from normal behavior of a cloud deployment that is not experiencing a ransomware attack, such that the mere presence of unusual activity may trigger the systems described herein to generate alerts or take some other action, even without explicit knowledge that the unusual activity is associated with a ransomware attack.

In some embodiments, particular policies may be put in place the systems described herein may be configured to enforce such policies as part of an effort to thwart ransomware attacks. For example, particular network sharing protocols (e.g., Common Internet File System ('CIFS'), Network File System ('NFS')) may be avoided when implementing storage for backup data, policies that protect backup systems may be implemented and enforced to ensure that usable backups are always available, multifactor authentication for particular accounts may be utilized and accounts may be configured with the minimum privilege required to function, isolated recovery environments may be created and isolation may be monitored and enforced to ensure the integrity of the recovery environment, and so on. As described in the present disclosure, the systems described herein may be configured to explicitly enforce such policies or may be configured to detect unusual activity that represents a violation of such policies, such that the mere presence of unusual activity may trigger the systems described herein to generate alerts or take some other action, even without explicit knowledge that the unusual activity is associated with a violation of a particular policy.

Readers will appreciate that ransomware attacks are often deployed as part of a larger attack that may involve, for example:

Penetration of the network through means such as, for example, stolen credentials and remote access malware.

Stealing of credentials for critical system accounts, including subverting critical administrative accounts that control systems such as backup, Active Directory ('AD'), DNS, storage admin consoles, and/or other key systems.

Attacks on a backup administration console to turn off or modify backup jobs, change retention policies, or even provide a roadmap to where sensitive application data is stored. Data theft attacks.

As a result of the many aspects that are part of a ransomware attack, embodiments of the present disclosure may be configured as follows:

The systems may include one or more components that detect malicious activity based on the behavior of a process.

The systems may include one or more components that store indicator of compromise ('IOC') or indicator of attack ('IOA') data for retrospective analysis.

The systems may include one or more components that detect and block fileless malware attacks.

The systems may include one or more components that remove malware automatically when detected.

The systems may include a cloud-based, SaaS-style, multi-tenant infrastructure.

The systems may include one or more components that identify changes made by malware and provide the recommended remediation steps or a rollback capability.

The systems may include one or more components that detect various application vulnerabilities and memory exploit techniques.

The systems may include one or more components that continue to collect suspicious event data even when a managed endpoint is outside of an organization's network.

The systems may include one or more components that perform static, on-demand malware detection scans of folders, drives, devices, or other entities.

The systems may include data loss prevention (DLP) functionality.

In some embodiments, the systems described herein may manage, analyze, or otherwise observe deployments that include deception technologies. Deception technologies allow for the use of decoys that may be generated based on

scans of true network areas and data. Such decoys may be deployed as mock networks running on the same infrastructure as the real networks, but when an intruder attempts to enter the real network, they are directed to the false network and security is immediately notified. Such technologies may be useful for detecting and stopping various types of cyber threats such as, for example, Advanced Persistent Threats ('APTs'), malware, ransomware, credential dumping, lateral movement and malicious insiders. To continue to outsmart increasingly sophisticated attackers, these solutions may continuously deploy, support, refresh and respond to deception alerts.

In some embodiments, the systems described herein may manage, analyze, or otherwise observe deployments that include various authentication technologies, such as multifactor authentication and role-based authentication. In fact, the authentication technologies may be included in the set of resources that are managed, analyzed, or otherwise observed as interactions with the authentication technologies may be monitored. Likewise, log files or other information retained by the authentication technologies may be gathered by one or more agents and used as input to the systems described herein.

In some embodiments, the systems described herein may be leveraged for the specific use case of detecting supply chain attacks. More specifically, the systems described herein may be used to monitor a deployment that includes software components, virtualized hardware components, and other components of an organization's supply chain such that interactions with an outside partner or provider with access to an organization's systems and data can be monitored. In such embodiments, supply chain attacks may be carried out in a manner such that patterns (e.g., specific interactions between internal and external systems) emerge, where the systems described herein may monitor for such patterns. Alternatively, supply chain attacks may involve behavior that deviates from normal behavior of a cloud deployment that is not experiencing a supply chain attack, such that the mere presence of unusual activity may trigger the systems described herein to generate alerts or take some other action, even without explicit knowledge that the unusual activity is associated with a supply chain attack.

In some embodiments, the systems described herein may be leveraged for other specific use cases such as, for example, detecting the presence of (or preventing infiltration from) cryptocurrency miners (e.g., bitcoin miners), token miners, hashing activity, non-fungible token activity, other viruses, other malware, and so on. As described in the present disclosure, the systems described herein may monitor for such threats using known patterns or by detecting unusual activity, such that the mere presence of unusual activity may trigger the systems described herein to generate alerts or take some other action, even without explicit knowledge that the unusual activity is associated with a particular type of threat, intrusion, vulnerability, and so on.

The systems described herein may also be leveraged for endpoint protection, such the systems described herein form all of or part of an endpoint protection platform. In such an embodiment, agents, sensors, or similar mechanisms may be deployed on or near managed endpoints such as computers, servers, virtualized hardware, internet of things ('IoT') devices, mobile devices, phones, tablets, watches, other personal digital devices, storage devices, thumb drives, secure data storage cards, or some other entity. In such an example, the endpoint protection platform may provide functionality such as:

Prevention and protection against security threats including malware that uses file-based and fileless exploits.

The ability to apply control (allow/block) to access of software, scripts, processes, microservices, and so on.

The ability to detect and prevent threats using behavioral analysis of device activity, application activity, user activity, and/or other data.

The ability for facilities to investigate incidents further and/or obtain guidance for remediation when exploits evade protection controls

The ability to collect and report on inventory, configuration and policy management of the endpoints.

The ability to manage and report on operating system security control status for the monitored endpoints.

The ability to scan systems for vulnerabilities and report/manage the installation of security patches.

The ability to report on internet, network and/or application activity to derive additional indications of potentially malicious activity.

Example embodiments are described in which policy enforcement, threat detection, or some other function is carried out by the systems described herein by detecting unusual activity, such that the mere presence of unusual activity may trigger the systems described herein to generate alerts or take some other action, even without explicit knowledge that the unusual activity is associated with a particular type of threat, intrusion, vulnerability, and so on. Although these examples are largely described in terms of identifying unusual activity, in these examples the systems described herein may be configured to learn what constitutes ‘normal activity’—where ‘normal activity’ is activity observed, modeled, or otherwise identified in the absence of a particular type of threat, intrusion, vulnerability, and so on. As such, detecting ‘unusual activity’ may alternatively be viewed as detecting a deviation from ‘normal activity’ such that ‘unusual activity’ does not need to be identified and sought out. Instead, deviations from ‘normal activity’ may be assumed to be ‘unusual activity’.

Readers will appreciate that while specific examples of the functionality that the systems described herein can provide are included in the present disclosure, such examples are not to be interpreted as limitations as to the functionality that the systems described herein can provide. Other functionality may be provided by the systems described herein, all of which are within the scope of the present disclosure. For the purposes of illustration and not as a limitation, additional examples can include governance, risk, and compliance (‘GRC’), threat detection and incident response, identity and access management, network and infrastructure security, data protection and privacy, identity and access management (‘IAM’), and many others.

In order to provide the functionality described above, the systems described herein or the deployments that are monitored by such systems may implement a variety of techniques. For example, the systems described herein or the deployments that are monitored by such systems may tag data and logs to provide meaning or context, persistent monitoring techniques may be used to monitor a deployment at all times and in real time, custom alerts may be generated based on rules, tags, and/or known baselines from one or more polygraphs, and so on.

Although examples are described above where data may be collected from one or more agents, in some embodiments other methods and mechanisms for obtaining data may be utilized. For example, some embodiments may utilize agent-less deployments where no agent (or similar mechanism) is deployed on one or more customer devices, deployed within

a customer’s cloud deployment, or deployed at another location that is external to the data platform. In such embodiments, the data platform may acquire data through one or more APIs such as the APIs that are available through various cloud services. For example, one or more APIs that enable a user to access data captured by Amazon CloudTrail may be utilized by the data platform to obtain data from a customer’s cloud deployment without the use of an agent that is deployed on the customer’s resources. In some embodiments, agents may be deployed as part of a data acquisition service or tool that does not utilize a customer’s resources or environment. In some embodiments, agents (deployed on a customer’s resources or elsewhere) and mechanisms in the data platform that can be used to obtain data from through one or more APIs such as the APIs that are available through various cloud services may be utilized. In some embodiments, one or more cloud services themselves may be configured to push data to some entity (deployed anywhere), which may or may not be an agent. In some embodiments, other data acquisition techniques may be utilized, including combinations and variations of the techniques described above, each of which is within the scope of the present disclosure.

Readers will appreciate that while specific examples of the cloud deployments that may be monitored, analyzed, or otherwise observed by the systems described herein have been provided, such examples are not to be interpreted as limitations as to the types of deployments that may be monitored, analyzed, or otherwise observed by the systems described herein. Other deployments may be monitored, analyzed, or otherwise observed by the systems described herein, all of which are within the scope of the present disclosure. For the purposes of illustration and not as a limitation, additional examples can include multi-cloud deployments, on-premises environments, hybrid cloud environments, sovereign cloud environments, heterogeneous environments, DevOps environments, DevSecOps environments, GitOps environments, quantum computing environments, data fabrics, composable applications, composable networks, decentralized applications, and many others.

Readers will appreciate that while specific examples of the types of data that may be collected, transformed, stored, and/or analyzed by the systems described herein have been provided, such examples are not to be interpreted as limitations as to the types of data that may be collected, transformed, stored, and/or analyzed by the systems described herein. Other types of data can include, for example, data collected from different tools (e.g., DevOps tools, DevSecOps, GitOps tools), different forms of network data (e.g., routing data, network translation data, message payload data, Wifi data, Bluetooth data, personal area networking data, payment device data, near field communication data, metadata describing interactions carried out over a network, and many others), data describing processes executing in a container, lambda, EC2 instance, virtual machine, or other execution environment), information describing the execution environment itself, and many other types of data.

Advantages and features of the present disclosure can be further described by the following statements:

1. A method comprising: accessing, by a data platform, a custom policy developed by a particular client entity that uses the data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store; determining, by the data platform, that the custom policy is likely to be of value to one or more other client

71

- entities, besides the particular client entity, that also use the data platform; and generating, by the data platform and in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities.
2. The method of any of the preceding statements, wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on an identity of the particular client entity.
 3. The method of any of the preceding statements, wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on a popularity metric of the custom policy exceeding a predetermined popularity threshold.
 4. The method of any of the preceding statements, further comprising: receiving, by the data platform from the particular client entity, user input representative of a request to publish the custom policy as a community policy; wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on the user input representative of the request to publish the custom policy.
 5. The method of any of the preceding statements, wherein: the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.
 6. The method of any of the preceding statements, further comprising: managing, by the data platform, a policy repository configured to store a plurality of community policies that are available for use by a plurality of different client entities; and storing, by the data platform and in response to the generating of the community policy, the community policy within the policy repository.
 7. The method of any of the preceding statements, wherein the managing of the policy repository includes: storing documentation corresponding to each community policy stored within the policy repository; and providing documentation corresponding to a particular community policy to a particular client device accessing the particular community policy from the policy repository.
 8. The method of any of the preceding statements, further comprising providing, by the data platform and for use by a user of the data platform, a graphical user interface that provides graphical assistance configured to assist a user in performing one or more of: discovering default policies and community policies available for use by a client entity with which the user is associated; creating new custom policies; or requesting to publish custom policies as community policies that are available for use by other users associated with other client entities besides the client entity with which the user is associated.

72

9. The method of any of the preceding statements, wherein the query invoked by the custom policy is a custom query that is also developed by the particular client entity.
10. The method of any of the preceding statements, wherein the query invoked by the custom policy is a default query provided by the data platform to a plurality of client entities including the particular client entity.
11. The method of any of the preceding statements, wherein the custom policy is at least one of: a compliance policy configured to verify that the target dataset complies with a predetermined compliance protocol; or a configuration policy configured to detect, within the target dataset, a particular attribute defined by the query invoked by the configuration policy.
12. The method of any of the preceding statements, further comprising constructing, by the data platform and based on client application data ingested from the cloud environment and loaded into the data store, a graph comprising a plurality of nodes connected by a plurality of edges, wherein each node of the plurality of nodes represents a logical entity from the client application data and each edge of the plurality of edges represents a behavioral relationship between nodes connected by the edge.
13. A system comprising: a memory storing instructions; and a processor communicatively coupled to the memory and configured to execute the instructions to: access a custom policy developed by a particular client entity that uses a data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store; determine that the custom policy is likely to be of value to one or more other client entities, besides the particular client entity, that also use the data platform; and generate, in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities.
14. The system of any of the preceding statements, wherein: the processor is further configured to execute the instructions to receive, from the particular client entity, user input representative of a request to publish the custom policy as a community policy; and the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on at least one of: an identity of the particular client entity, a popularity metric of the custom policy exceeding a predetermined popularity threshold, or the user input representative of the request to publish the custom policy.
15. The system of any of the preceding statements, wherein: the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.
16. The system of any of the preceding statements, wherein the processor is further configured to execute the instructions to: manage a policy repository config-

ured to store a plurality of community policies that are available for use by a plurality of different client entities; and store, in response to the generating of the community policy, the community policy within the policy repository.

17. A computer program product embodied in a non-transitory computer readable storage medium and comprising computer instructions for: accessing a custom policy developed by a particular client entity that uses a data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store; determining that the custom policy is likely to be of value to one or more other client entities, besides the particular client entity, that also use the data platform; and generating, in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities.
18. The computer program product of any of the preceding statements, further comprising computer instructions for receiving, from the particular client entity, user input representative of a request to publish the custom policy as a community policy; wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on at least one of: an identity of the particular client entity, a popularity metric of the custom policy exceeding a predetermined popularity threshold, or the user input representative of the request to publish the custom policy.
19. The computer program product of any of the preceding statements, wherein: the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.
20. The computer program product of any of the preceding statements, further comprising computer instructions for: managing a policy repository configured to store a plurality of community policies that are available for use by a plurality of different client entities; and storing, in response to the generating of the community policy, the community policy within the policy repository.

Various processes and techniques described herein relate to creating a contextual model of a large dataset (e.g., a graph or mapping of the data a client entity stores in cloud environments described herein) and using that contextual model to identify anomalies, security threats, compliance issues, and so forth. For these types of objectives, a behavior-based approach (e.g., in which behaviors are monitored and analyzed to identify anomalies in light of a contextual understanding of the dataset, etc.) allows data platforms to discover and identify complex issues in various ways that have been described. For example, such approaches may assist a client entity (e.g., a company, organization, customer, user, group of users, etc., that uses, subscribes to, relies on, benefits from, and/or otherwise is a client of data platforms described herein) in becoming efficiently and

effectively apprised of various types of complex issues. This may be beneficial at least in part because these complex issues require a deep understanding of data context to identify, and hence may tend to be difficult to search for and uncover for human experts and computing systems alike.

Along with behavior-based approaches such as have been described, many successful data management strategies (e.g., security strategies deployed by client entities, etc.) may also leverage policy-based approaches that assist the client entity in becoming efficiently and effectively apprised of various types of issues that may not necessarily require the same contextual understanding. For example, certain issues may be searched for and identified based on known patterns and actions that may be well-defined and understood by the client entity. For these types of issues, a policy-based approach (e.g., in which policies and queries are defined to handle specific, well-defined scenarios in particular ways) may complement the behavior-based approach described above to maintain secure and well-managed data.

One way that a data platform may provide or support such a policy-based approach is to provide standardized or default policies (e.g., AWS, Azure, GCP, or other host policies) for use by client entities on the data platform. Many client entities may benefit from this type of one-size-fits-all approach, at least for a time or while the datasets they are managing are relatively small and/or conventional. As a client entity's data needs grow and evolve over time, however, the client entity may find that default policies increasingly fail to perfectly fit their needs. In part to address these types of client needs, data platforms described herein may therefore provide ways that client entities may edit and modify default policies, as well as build, update, validate, share, and otherwise create new policies. Any of these types of custom policies may be configured to customize policy-based strategies implemented by the data platform to specific needs of a particular client entity.

While one way to support such customization of data policies may be to allow client entities to create requests for default policies to be modified or for new default policies to be created (e.g., by the provider of the data platform), this type of customization approach may be associated with various costs and drawbacks that may make it non-ideal. For example, this type of approach may not scale well as a data platform grows and a large number of client entities (all with different security and data management goals, strategies, needs, best practices, etc.) all desire different types of customizations.

Accordingly, as will now be described with reference to FIGS. 5A-7C, data platforms described herein may be configured to provide or support standardized and/or proprietary query languages, policy-creation tools, structured application programming interfaces (APIs), user interfaces (UIs) (including graphical user interfaces (GUIs)), and/or other tools and implements to facilitate the creation of custom policies and queries that cater to the needs of the variety of different client entities that may use the data platforms.

More specifically, as described in detail below, FIGS. 5A-5K illustrate how a query language associated with a data platform (e.g., a proprietary query language associated with any of the data platforms described herein, a standardized query language such as SQL, etc.) may be used to define queries used to search for and identify unwanted behavior in particular datasets or data environments. The query language may have a human-readable text syntax

75

configured to facilitate specifying the selection, filtering, and manipulation of data, and may be convertible to other languages including standardized query languages such as SQL. FIGS. 5A-5K also illustrate how custom policies (which may be written in the same query language or using associated syntax that is readily processed by the query language) may invoke queries such as described above and may annotate query results to, for example, raise alerts, direct reports to be generated, or otherwise manage the visibility that users will have into the queried data. FIGS. 6A-6E then illustrate additional details about how alert profiles may be used in connection with custom policies and queries, and how alert profiles may be modular so as to be invoked by more than one custom policy. FIGS. 7A-7C illustrate how useful custom policies and/or queries may be shared and distributed between client entities and how the data platform may generate or designate custom policies to be what are referred to as “community policies” that are available for use by multiple client entities.

As will be made apparent, various benefits may arise from custom policies and queries, modular alert profiles, community policies, and various other principles and concepts described in connection with the following figures. For example, while default rules and policies may not always match a client entity’s needs, the ability to customize policies and queries to the known challenges and conditions that a particular client entity faces may be a significant advantage. Moreover, using custom (e.g., proprietary or other) query languages in the ways described herein, client entities may be enabled express security objectives as code, which may be desirable in light of evolving paradigms such as Policy as Code, Infrastructure as Code, declarative infrastructure models, and so forth. Another benefit of the principles described below involves the separation of queries written in the query language and the custom policies that invoke those queries. As will be made apparent, the paradigm described below emphasizes modularity and reusability such that policy owners may be different entities from query writers (e.g., coders working in the query language, etc.) and may attach queries (e.g., from a library or other repository) to custom policies without necessarily needing to understand all the implementation details of the queries.

FIG. 5A illustrates an example of a process 500 for implementing custom policies and queries for a data platform. While FIG. 5A shows illustrative operations according to one implementation, other implementations may omit, add to, reorder, and/or modify any of the operations shown in FIG. 5A. In some examples, multiple operations shown in FIG. 5A or described in relation to FIG. 5A may be performed concurrently (e.g., in parallel) with one another, rather than being performed sequentially as illustrated and/or described. One or more of the operations shown in FIG. 5A may be performed by a data platform described herein (e.g., data platform 12). More particularly, computing resources of a data platform (e.g., data ingestion resources 18, data processing resources 20, etc.) may perform each of the operations of process 500 during or after a data ingestion procedure has been performed to move a target dataset from a cloud environment (e.g., cloud environment 14) to a data store (e.g., data store 30). The operations of process 500 may be performed prior to, in parallel with, or subsequent to other operations described above to be performed by the data platform. For example, operations 501-503 may be performed before, during, or after the data platform constructs a graph (e.g., a polygraph) in any of the ways described herein. Based on client application data ingested from the cloud environment (e.g., cloud environment 14) and loaded

76

into the data store (e.g., data store 30), the data platform (e.g., data platform 12) may construct a graph (e.g., a polygraph) comprising a plurality of nodes connected by a plurality of edges. As described above, each node of the plurality of nodes may represent a logical entity from the client application data and each edge of the plurality of edges may represent a behavioral relationship between nodes connected by the edge.

Each of operations 501-503 of process 500 will now be described in more detail as the operations may be performed by computing resources of a data platform described herein. The graph construction and/or other such operations described above may be performed in connection with operations 501-503 in any suitable way.

At operation 501, the data platform (e.g., computing resources operating on the data platform) may access a custom policy. As will be described and illustrated in more detail below, the custom policy may invoke a query that is written in a query language associated with the data platform. For example, as mentioned above, the query language may be a standardized query language (e.g., SQL) or a proprietary query language that is specifically configured to operate (or otherwise associate with) the data platform. The query invoked by the custom policy accessed at operation 501 may itself be a custom query (e.g., drafted by a same person or same client entity as the custom policy, drafted by a different person or different client entity as the custom policy, etc.) in certain examples. In other examples, the query may be a default query selected from a library of queries provided by the data platform to be invoked by custom policies. While this example and various other examples described herein refer to a (singular) query being invoked by the custom policy, it will be understood that, in certain implementations, a custom policy (or other policies such as default or community policies described herein) may invoke a plurality of queries (e.g., custom queries, default queries, community queries, a combination of these, etc.) that may each serve different purposes and that may inter-operate to advance the objectives of the policy in any manner as may serve a particular implementation. In any case, the query (or each of the queries, in examples in which the custom policy at operation 501 invokes more than one) may define an attribute that is to be identified within a target dataset that is being (or has been) ingested from the cloud environment to be stored in the datastore. The custom policy may define an action associated with the attribute, such as an alert or other notification that is to be delivered in the event that the attribute is identified within the target dataset. Examples of custom policies and queries will be provided below.

At operation 502, the data platform may detect, based on the query and within the target dataset, data characterized by the attribute. For instance, as will be described in more detail in an extended example below, the target dataset could include trail data from the cloud environment and the query could define an attribute of the trail data that indicates a sign-in attempt from an unknown IP address (e.g., an IP address not included on a predefined list of expected addresses). In this example, the data platform may therefore, at operation 502, detect within the trail data that the sign-in attempt occurred.

At operation 503, the data platform may then provide a query result associated with the data characterized by the attribute. For example, the query result may be provided (e.g., by way of an API) to a user interface being viewed by a user of the data platform (e.g., an owner or other person responsible for the target dataset). The provided query result

may be annotated in accordance with the custom policy to indicate any suitable information. For example, the query result may be annotated to indicate the detecting of the data characterized by the attribute (e.g., to indicate that the sign-in attempt from the unknown IP address has been detected) and to direct the action to be performed (e.g., to alert the user in accordance with an alert defined by or invoked by the custom policy, to recommend a remedial or other action be performed by the user, etc.).

To illustrate the principles introduced by process 500 in more detail, FIG. 5B shows an example policy that invokes a query and an alert profile to provide an annotated query result to a user. Specifically, as shown, a policy 510 may invoke a query 511 (which defines an attribute) and an alert profile 512. Query 511 is shown to detect data characterized by the attribute within a target dataset 513, and to provide a query result 514, by way of an API 515, to a user interface 516 being used by a user 517. As described above in relation to process 500, query result 514 may be annotated, in accordance with policy 510, with annotations 518 that include indications (e.g., an indication of the detecting of the data characterized by the attribute, etc.) and/or direction (e.g., directing an action defined by policy 510 to be performed by user 517, etc.). Each of the elements of FIG. 5B will now be described in more detail.

Policy 510 may be implemented by any mechanism used to add annotated metadata to a query (e.g., query 511) for improving the context of alerts, reports, and/or other information displayed in a user interface such as user interface 516. Policy 510 may be used to monitor behavior in target dataset 513 such as a set of trail data (e.g., AWS CloudTrail data or the like), and may be managed by way of API 515 using a query language associated with the data platform. Once policy 510 is created, it may be presented in a “new policies” portion of user interface 516 so as to be viewed, used, and managed by user 517.

Policy 510 may be implemented as any of various types of policies such as a compliance policy configured to verify that target dataset 513 complies with a predetermined compliance protocol, a configuration policy configured to detect (within target dataset 513) the particular attribute defined by the query 511 invoked by the configuration policy, or another suitable type of policy. A “policy_type” tag in policy 510 or query 511 may be used to indicate the type of policy (or query) that is being implemented in a given implementation (e.g., trail data activity policies vs. compliance policies, event vs. CTA queries, etc.). Different types of policies 510 and/or queries 511 may provide and empower client entities to create custom rules catering to their security posture, may provide astute methods to discover compliance and configuration discrepancies for cloud resources, may enable customers to add missing policies for existing reports, and may achieve other such benefits and advantages. Client entities using the data platform may desire to create their own compliance and configuration policies to appropriately build up the inventory suitable for their specific audit needs. Custom policies, custom queries, custom alert profiles, and so forth, all may help enable these client entities to create alerts and notifications based on their own understanding of their data environments.

Query 511 may be implemented by any mechanism used to interactively request information from a specific curated dataset (e.g., target dataset 513). Queries such as query 511 may have defined structures for authoring detections. For example, query 511 may be authored in a particular query language (e.g., a standard or proprietary query language) that is associated with the data platform. As will be illus-

trated in more detail in examples provided below, query 511 may define various aspects of an attribute (“Attribute” in FIG. 5B) that is to be detected, as well as certain other details such as which dataset should be searched to detect the attribute, what information is to be exposed when the attribute is detected, and so forth. Specifically, for instance, query 511 may define a source dataset where the detections are to be made (e.g., “CloudTrailRawEvents” to define a dataset with raw AWS CloudTrail data that is ingested from an AWS cloud environment, etc.); a filter specifying records of interest that make up the attribute being searched for (e.g., “EVENT_SOURCE=‘iam.amazonaws.com’ and EVENT_NAME=‘DeleteAccessKey’ and ERROR_CODE is null” for a particular attribute associated with detecting that an access key is deleted, etc.); and certain fields that this query is to expose when the attribute is detected (e.g., a “return distinct” clause indicating that unwanted duplicates should be ignored and specifying fields such as “INSERT_ID, INSERT_TIME, EVENT_TIME, EVENT”, etc.). Example code for a query such as query 511 will be illustrated and described in more detail below.

The attribute defined by query 511 may describe a particular anomaly, a behavior of interest, or another type of attribute. In certain examples, the attribute characterizing the data detected based on query 511 may correspond to a predetermined data security threat (e.g., a risky behavior that has been determined to correlate with security issues, etc.). In other examples, the attribute characterizing the data detected based on query 511 may correspond to a failure to satisfy a predetermined compliance rule. In any case, when the attribute defined in query 511 can be detected (e.g., identified) within target dataset 513, query 511 may return query result 514 that is reported and handled in accordance with policy 510.

In certain examples, policy 510 and/or query 511 may be provided by the data platform to serve as a default policy/query for use by various client entities. In other examples, policy 510 and/or query 511 may be a custom policy/query that is fully customized and/or customer-defined (e.g., authored by user 517 or another person associated with the client entity of user 517). To illustrate, FIG. 5C shows examples of different types of policies invoking different types of queries. Specifically, as shown, the data platform may provide both a default policy 510-D1 and a default query 511-D1 that default policy 510-D1 invokes in certain examples. In other examples, a client entity may author (and fully customize) both a custom policy 510-C1 and a custom query 511-C1 that custom policy 510-C1 invokes. Custom policies and/or queries may be developed from scratch or, in certain circumstances, may be generated based on (e.g., by modifying) default policies and/or queries that already exist and are provided by the data platform (e.g., default policy 510-D1 and/or default query 511-D1). More particularly, for instance, custom policy 510-C1 may be generated based on a default policy (e.g., default policy 510-D1) that is provided by the data platform such that custom policy 510-C1 inherits one or more attributes of the default policy while also modifying one or more other attributes of the default policy. Similarly, custom query 511-C1 may be generated based on a default query (e.g., default query 511-D1) that is provided by the data platform such that custom query 511-C1 inherits one or more attributes of the default query while also modifying one or more other attributes of the default query. As further shown in FIG. 5C, in certain examples, a default policy 510-D2 may invoke a custom query 511-C2, or a custom policy 510-C2 may invoke a default query 511-D2.

Returning to FIG. 5B, alert profile **512** represents instructions that may be accessed or invoked by policy **510** to determine how to handle alerts (e.g., system notifications based on violation of configured policies, etc.) that may be triggered by query **511** identifying the attribute in target dataset **513**. Alerts may be presented on an event triage dossier or on an event timeline widget, and may be processed/filtered by an existing alert routing component. Any of these or other such operations that are to be taken with respect to alerts may be defined by alert profile **512**. As will be described in more detail below, certain alert profiles may be modular so as to be invoked by different custom policies authored by different client entities.

Target dataset **513** may be implemented by any suitable dataset in which a particular attribute is to be detected. As one example mentioned above, for instance, target dataset **513** may be implemented by trail data representative of events (e.g., raw events) that have been recorded as having occurred in a cloud environment. If the cloud environment is associated with Amazon Web Services (AWS), for example, such trail data may be CloudTrail data provided by AWS. Similarly, if the cloud environment is associated with Google Cloud Platform (GCP), the trail data may be Cloud Audit Logs provided by GCP; if the cloud environment is associated with Microsoft Azure, the trail data may be Log Analytics provided by Azure; and so forth for whichever cloud environment may be used. In other examples, rather than trail data representing events that have occurred with respect to client application data stored in the cloud environment, target dataset **513** may instead be implemented by the client application data itself, or any other source data or metadata as a client entity may desire to query in a particular implementation.

Based on the detecting of the attribute by query **511**, query result **514** may be generated in accordance with policy **510** and/or alert profile **512** and may be provided by way of API **515** for presentation by user interface **516** to user **517**. As shown, query result **514** may include the various annotations **518** defined by policy **510** (as well as, in certain examples, defined by query **511** and/or alert profile **512**). Specifically, for instance, query result **514** may be annotated in accordance with policy **510** to indicate the detecting of the data characterized by the attribute and to direct an action to be performed. Based on policy **510** and/or alert profile **512**, the directing of the action to be performed may include directing user interface **516** of the data platform to display an alert or a report to user **517**. For instance, the alert or the report may indicate the detecting of the data characterized by the attribute, may explain the nature of the attribute (e.g., why it is an issue if it represents a security threat or compliance failure, etc.), and/or may provide remedial or other recommended actions to be taken by the user (e.g., to remedy, mitigate, or otherwise respond appropriately to the detection of the attribute). In some examples, the directing of the action to be performed may involve the data platform directing performance of an automatic operation responsive to query result **514** (e.g., in addition or as an alternative to alerting user **517** and/or directing them to perform the remedial action). For instance, the data platform may direct one or more agents operating within the cloud environment (e.g., agents **38** or other agents described herein) to terminate or to behave differently as they push data to the data platform for ingestion. As another example, the data platform may alter the way a graph (e.g., a polygraph) is being generated, or otherwise adjust an operation described herein or perform another suitable operation to help address or mitigate the conditions represented by the detected attribute.

API **515** may represent any suitable application programming interface whereby computing resources executing query **511** (as well as annotating query result **514** in accordance with policy **510** and/or alert profile **512**) may communicate with user interface **516** to apprise user **517** of that query result. To this end, the data platform may provide an API **515** that supports the query language in which query **511** is authored, and that is configured to facilitate query and custom policy creation in any suitable way. For example, the accessing of a custom policy (e.g., policy **510**) and the providing of a query result (e.g., query result **514**), as described above in relation to process **500**, may both be performed by way of API **515**. Moreover, along with facilitating the creating of custom policies and/or queries, API **515** may further be configured to facilitate accessing, updating (e.g., editing, modifying, etc.), validating, viewing (e.g., listing, reviewing details of, etc.), sharing, publishing, deleting, and/or otherwise managing queries and custom policies.

User **517** may represent any person associated with a particular client entity (e.g., company, organization, or other client/customer of the data platform) that is responsible for creating and/or executing policy **510** and query **511**. For instance, user **517** may be an employee of a company and may be responsible for overseeing data security and/or compliance of that company's data resources.

User interface **516** may represent any suitable user interface (e.g., graphical user interface, command-line interface, etc.) that user **517** may interact with to manage policies, queries, alert profiles, and so forth, and to gain visibility into query results such as query result **514**. For example, the data platform may provide, for use by user **517**, a graphical user interface (GUI) that implements API **515** and provides graphical assistance to user **517** as user **517** creates at least one of query **511** or policy **510** (e.g., as a fully or partially customized query or policy). As will be described in more detail below, certain custom policies and/or custom queries may be converted into or used as a basis for generating policies and/or queries for more general use (referred to herein as community policies or queries), similar to default policies and/or queries. In such cases, user interface **516** may represent the providing, by the data platform for use by user **517**, of a GUI that provides graphical assistance configured to assist user **517** in performing one or more of 1) discovering default policies and community policies available for use by a client entity with which user **517** is associated; 2) creating new custom policies; or 3) requesting to publish custom policies as community policies (or otherwise share custom policies) to be available for use by other users associated with other client entities besides the client entity with which **517** user is associated. Examples of community policies will be provided with additional description below.

Using user interface **516**, user **517** may access any functions implemented by API **515**. For example, user **517** may manage custom policies, queries, and/or alert profiles (e.g., creating, editing, deleting, etc.); make use of default and/or community policies, queries, and/or alert profiles (e.g., accessing these in a read-only manner, etc.); receive and/or provide direction as to how certain query results are to be addressed (e.g., receiving remedial action direction as an annotation **518** of a query result **514**, directing the data platform to perform certain operations, etc.); and so forth.

The various operations associated with process **500**, as laid out in FIG. 5A and as illustrated and described further in relation to FIGS. 5B and 5C, may be performed as part of or subsequent to a data ingestion process whereby data is ingested from a cloud environment and stored in a data store

for further processing by a data platform. For instance, in certain implementations, the data platform may perform a data ingestion process (whereby client application data and/or the target dataset are loaded from a cloud platform into a data store for processing by the data platform) and the detecting of the data characterized by the attribute (operation **502**) and the providing of the query result (operation **503**) may be performed during the data ingestion process prior to the target dataset being completely loaded into the data store. In contrast, in other implementations (or under other conditions in the same implementations), the data platform may perform the data ingestion process (whereby the client application data and/or the target dataset are likewise loaded from the cloud platform into the data store for processing by the data platform) and the detecting of the data characterized by the attribute (operation **502**) and the providing of the query result (operation **503**) may be performed subsequent to the data ingestion process when the target dataset is completely loaded into the data store.

To illustrate, FIG. 5D shows an example configuration in which a data platform **520** may query different types of target datasets at different stages of the data ingestion (e.g., during the data ingestion, after the data ingestion, etc.). Specifically, as shown in FIG. 5D, data ingestion resources **521** of data platform **520** may be used to ingest (e.g., by either push or pull techniques, as have been described herein) cloud data **522** from a cloud environment **523** (e.g., representing any cloud environment described herein) to a data store **524** (e.g., representing any data store described herein). Within data store **524**, FIG. 5D shows two potential embodiments of target dataset **513**: 1) a target dataset **513-1** implemented by a set of trail data **525** pulled from cloud environment **523** in any of the ways described herein and comprising various event records (e.g., cloud activity events exposed in raw format, etc.); and 2) a target dataset **513-2** implemented by a set of client application data **526** comprising actual application data (e.g., control plane data) that may be pushed by agents in cloud environment **523** in any of the ways described herein. Data processing resources **527** may represent any of the data processing resources of a data platform described herein, and, as described in detail above, may be used, for example to analyze target datasets within data store **524** for any suitable purpose (e.g., generating a polygraph, understanding data context, identifying anomalies or other issues, etc.).

Data ingestion resources **521** may perform operations of process **500** with respect to either or both target dataset **513** implementations in any way and at any point in the process as may serve a particular implementation. For example, in certain implementations or under certain circumstances, data ingestion resources **521** may analyze cloud data **522** in real time during the ingestion process (e.g., before trail data **525** and/or client application data **526** are fully loaded into data store **524**) so as to identify attributes of interest as early in the process as possible. In other implementations or under other circumstances, however, data ingestion resources **521** may complete the ingestion of cloud data **522** such that trail data **525** and/or client application data **526** are fully loaded into data store **524**, and these target datasets may then be queried (e.g., by data ingestion resources **521** or data processing resources **527**) from data store **524**.

As illustrated by trail data **525**, a particular target dataset **513-1** may comprise trail data received from a cloud environment (e.g., cloud environment **523**) that stores and executes client application data (e.g., client application data **526**). As shown, this trail data **525** may include a plurality of event records generated within the cloud environment

with respect to the client application data. The data detected to be characterized by the attribute within this target dataset (at operation **502** of process **500**) may thus include one or more event records of the plurality of event records. As mentioned above, if this cloud environment (e.g., cloud environment **523**) is associated with AWS, the trail data (trail data **525**) may be implemented by CloudTrail data provided by AWS, whereas if the cloud environment is associated with GCP, Azure, or another cloud provider entity, the trail data may be a different, but analogous, collection of data (e.g., cloud audit log data, log analytics data, etc.).

FIG. 5E illustrates an example of a query **531** and a custom policy **532** that invokes query **531**. In this example, the attribute that the data platform is attempting to detect is the presence of a service account login from an unknown IP address. More particularly, one or more of the event records of trail data **525** may be determined to represent an attempt (or multiple attempts) to log into a service account of cloud environment **523** from one or more unauthorized IP addresses (e.g., addresses not included on a specific “whitelist” of approved or recognized IP addresses). For purposes of illustration, for example, the IP addresses “1.1.1.1” and “2.2.2.2” may be considered to be known or approved addresses from which to log in, while all other IP addresses may be considered unknown.

As shown in query **531**, a query entitled “Global_AWS_CTA_LoginUnknownIP” may be generated using appropriate syntax of the query language that is associated with the data platform. Query **531** is shown to define a data source (“source”) as AWS trail data events (“CloudTrailRawEvents”), to define a sought-after attribute (“filter”) as a sign-in event (an EVENT_SOURCE of “signin.amazonaws.com” and an EVENT: useridentity.type of “AWSService”) in which a sourceIPAddress is neither ‘1.1.1.1’ or ‘2.2.2.2’, and where the ERROR_CODE is null. For trail data events detected to be characterized by such attributes, query **531** designates that fields INSERT_ID, INSERT_TIME, EVENT_TIME, and EVENT are to be exposed.

As shown in policy **532**, query **531** may be invoked by setting “queryId” to “Global_AWS_CTA_LoginUnknownIP”, and various other parameters configured to manage the annotated metadata added to the query results may also be set. Specifically, these parameters are shown in FIG. 5E to include examples such as a “policyType” set to “Violation”, a “title” set to “Service Account Login from Unknown IP Address”, an “enabled” flag set to true, a policy “description” set to “A service account logged in from an unexpected IP address”, a “remediation” direction set to “Investigate any suspicious activity”, a “severity” of identifying this attribute set to “high”, a “limit” of attributes to return set to 100, a time frequency “evalFrequency” at which the query is to be executed set to “Hourly”, an “alertEnabled” flag set to true, an invocation of a particular “alertProfile” set to “LW_CloudTrail_Alerts”, a “policyUi” including a “domain” set to “AWS” and a “subdomain” set to “Cloudtrail”, and an “evaluatorId” set to “Cloudtrail”. Accordingly, in this example, a query response may expose an “INSERT_ID”, an “INSERT_TIME”, an “EVENT_TIME,” and various details of an EVENT that matches the defined attribute.

While in certain examples, query **531** and/or policy **532** may each be customized and authored “from scratch” (i.e., without being based on any other existing query or policy, such as a default query or policy), it is also possible, in other examples, that either or both of query **531** and policy **532**

may be customized by updating or modifying an existing query or policy. For example, a default policy may exist that is identical to custom policy 532 except that the default policy sets “severity” to “low” instead of to “high”. In this example, custom policy 532, instead of being authored from scratch, may be generated by using the default policy as a starting point and overriding the “severity” parameter to set the severity to “high” (rather than “low”) to better suit the needs of the client entity creating the custom policy. Other parameters of either query 531 or custom policy 532 may be similarly altered or updated so as to leverage default or other custom queries and policies as a starting point in generating new custom queries and/or policies.

While query 531 and custom policy 532 provide one example of a custom policy invoking a query to thereby provide a benefit to a user of the data platform (e.g., in this case, to find instances of when someone with an unknown IP address attempted to log in and access a restricted account), it will be understood that this policy and query are illustrative only and that various other custom policies and/or queries may be generated and used in other examples and/or for other purposes. To provide a few additional examples, FIGS. 5F-5K each illustrate another example of an illustrative query and custom policy invoking the query. Each of these examples will be understood to function similarly to the example of FIG. 5E described in detail above, though certain details may be specific to each example, as will now be described.

In FIG. 5F, an illustrative query 533 is shown to have the title: “LW_Custom AWS CTA_AuroraPasswordChange”. Query 533 is configured to view CloudTrail events to detect instances where the password for an Aurora RDS cluster has been changed. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 533), the data platform may provide a query result in accordance with a custom policy 534 having a policy ID of “CustomPolicy1” and invoking query 533 (i.e., “queryId: LW_Custom_AWS_CTA_AuroraPasswordChange”). As shown, custom policy 534 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user to take the remedial action of “Check[ing] that the password change was expected and ensure only specified users can modify the RDS cluster.”

In FIG. 5G, an illustrative query 535 is shown to have the title: “LW_Custom AWS CTA_RDSBackupDisabled”. Query 535 is configured to view CloudTrail events to detect instances where the backups have been disabled for an RDS database instance. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 535), the data platform may provide a query result in accordance with a custom policy 536 having a policy ID of “CustomPolicy2” and invoking query 535 (i.e., “queryId: LW_Custom_AWS_CTA_RDSBackupDisabled”). As shown, custom policy 536 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user to take the remedial action of “Ensuring only specified users can modify DB instances.” The action may further include providing direction to the user regarding how to reenabling the RDS backups: “To reenabling, navigate to the database, select Modify, set Backup Retention Period greater than 0.”

In FIG. 5H, an illustrative query 537 is shown to have the title: “LW_Custom_AWS CTA_RootAccountMFADisabled”. Query 537 is configured to view CloudTrail events to detect instances where the AWS root account multi-factor authentication (MFA) device was deactivated or deleted. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 537), the data platform may provide a query result in accordance with a custom policy 538 having a policy ID of “CustomPolicy3” and invoking query 537 (i.e., “queryId: LW_Custom_AWS_CTA_RootAccountMFADisabled”). As shown, custom policy 538 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user to take the remedial action of “Enabl[ing] MFA for the root account.” The action may further include specific instructions for how to do so (certain of the following steps are represented by “/* ADDITIONAL STEPS */” and not explicitly shown in FIG. 5H due to space constraints): “1. Log in to your account using the root credentials. 2. Choose the account name near the top-right corner of the page and then choose My Security Credentials. 3. Choose Multi-factor authentication (MFA). 4. Choose Activate MFA. 5. Choose the type of device to use for MFA and then choose Continue. 6. Complete the steps to configure the device type appropriate to your selection.”

In FIG. 5I, an illustrative query 539 is shown to have the title: “LW_Custom AWS CTA_KMSChange”. Query 539 is configured to view CloudTrail events to detect instances where an AWS Key Management Service (KMS) change has occurred. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 539), the data platform may provide a query result in accordance with a custom policy 540 having a policy ID of “CustomPolicy4” and invoking query 539 (i.e., “queryId: LW_Custom_AWS_CTA_KMSChange”). As shown, custom policy 540 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user to take the remedial action of “Check[ing] the KMS change and ensure only specified users can make KMS changes.”

In FIG. 5J, an illustrative query 541 is shown to have the title: “LW_Custom_AWS CTA_FailedLogin”. Query 541 is configured to view CloudTrail events to detect instances where a failed AWS console login attempt has occurred. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 541), the data platform may provide a query result in accordance with a custom policy 542 having a policy ID of “CustomPolicy5” and invoking query 541 (i.e., “queryId: LW_Custom_AWS_CTA_FailedLogin”). As shown, custom policy 542 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user to take the remedial action of “Investigat[ing] any suspicious activity.”

In FIG. 5K, an illustrative query 543 is shown to have the title: “LW_Custom AWS CTA_UnencryptedVolumeCreated”. Query 543 is configured to view CloudTrail events to detect instances where a compliance rule or best practice to encrypt EBS volumes has not been complied with. When such events are detected (i.e., when data within the CloudTrail dataset is found to be characterized by the attribute defined by the “filter” of query 543), the data platform may provide a query result in accordance with a custom policy

85

544 having a policy ID of “CustomPolicy6” and invoking query 543 (i.e., “queryId: LW_Custom_AWS_CTA UnencryptedVolumeCreated”). As shown, custom policy 544 may direct that the query result be annotated to indicate that the event was detected and to direct an action to be performed. Specifically, in this case, the action may be to advise the user that “It is not possible to encrypt an unencrypted EBS volume”, and to then “suggest creating a snapshot, then creating an encrypted volume with the snapshot” (or providing other such direction to help bring the data into compliance).

Various other examples not explicitly illustrated in the figures may include queries and custom policies facilitating detecting of security, configuration, compliance, and/or other relevant issues. To provide a few additional, non-limiting examples:

One custom policy and query could be created to detect and alert the user when the password for an RDS cluster has been scheduled for change and the password will be changed in the next maintenance window. A recommended remedial action in this case could be to check that the password change was expected and ensure only specified users can modify the RDS cluster.

Another custom policy and query could be created to detect and alert the user when a CloudWatch destination is deleted. A recommended remedial action in this case could be to ensure only specified users can delete CloudWatch destinations and possibly to advise the user to create a new CloudWatch destination.

Another custom policy and query could be created to detect and alert the user when CloudTrail (or other trail data) activity is detected for an IAM user on *Kali* Linux. A recommended remedial action in this case could be to investigate any suspicious activity.

Another custom policy and query could be created to detect and alert the user when a failed AWS console login attempt has been detected. A recommended remedial action in this case could be to investigate any suspicious activity.

Another custom policy and query could be created to detect and alert the user when a route 53 hosted zone has been deleted. A recommended remedial action in this case could be to ensure only specified users can delete hosted zones and possibly to advise the user to create a new hosted zone.

Another custom policy and query could be created to detect and alert the user when a rule has been added to a security group that allows unrestricted ingress. A recommended remedial action in this case could be to remove any inbound rules on security groups that allow unrestricted ingress.

Another custom policy and query could be created to detect and alert the user when a public ACL has been attached, or “AllUsers” or “AuthenticatedUsers” have been granted “FULL_CONTROL,” “READ,” “WRITE,” “READ_ACP,” or “WRITE_ACP” access. A recommended remedial action in this case could be to consider changing the access to this bucket with an ACL or policy so only intended accounts, users, or services have access.

Another custom policy and query could be created to detect and alert the user when an S3 bucket with “public-read,” “public-read-write,” or “authenticated-read” ACL has been created. A recommended remedial action in this case could be to advise the user that AWS does not recommend “public-read-write” access on buckets, and that the user may consider attaching an ACL or policy so only intended accounts, users, or services have access.

86

Another custom policy and query could be created to detect and alert the user when a policy allowing global read or write access has been attached to an S3 bucket. A recommended remedial action in this case could be to advise the user to consider changing the policy so only intended accounts, users, or service will have access.

Another custom policy and query could be created to detect and alert the user when versioning has been enabled for an S3 bucket (since rates apply for every version of an object stored or requested). A recommended remedial action in this case could be to advise the user that if versioning is not required, it can be suspended to save costs.

Another custom policy and query could be created to detect and alert the user when a policy allowing global write access has been attached to an S3 bucket. A recommended remedial action in this case could be to advise the user to consider changing the policy so only intended accounts, users, or services will have write access.

Various other custom policies and/or queries may similarly be implemented as may serve a particular implementation. Indeed, the flexibility of the query language and overall policy/query customization framework described herein to allow users to create and edit custom policies and queries to meet their own specific needs is a major benefit and advantage of the framework, as has been described.

Referring back to FIG. 5B, various implementations of policy 510 (including custom policies created from scratch or based on existing policies) have been described to invoke various implementations of query 511 to detect attributes of target dataset 513 and to annotate query results 514 in helpful ways (e.g., with alerts, notifications, indications, directions or instructions, etc.). In all of these examples, client entities making use of the data platform providing this policy/query framework (e.g., represented by a user 517 in FIG. 5B) may benefit from being able to express their particular security objectives, compliance rules, best practices, and so forth, as code using a query language associated with the data platform. As has been described and made apparent, this ability to codify various aspects of their data management strategy may provide client entities with great control, flexibility, and other benefits.

Also shown in FIG. 5B (and described briefly above) is an alert profile 512 that policy 510 may reference or invoke in a similar way as it invokes query 511. Alert profiles (such as the alert profile 512 shown in FIG. 5B) may serve as a contract between multiple components or microservices managed and implemented by the data platform. For example, an alert profile may define contract provisions between a policy and an event generator service (described in more detail below) or may set forth constraints to be enforced by the event generator service. As will be shown, a drafter of a policy may associate the policy with an instance of the alert profile to thereby invoke this contract and exert control over how alerts stemming from the policy will be treated by the event generator service and other microservices (e.g., a triage service prioritizing alerts, etc.). The definition of an alert profile may follow a typed schema design (i.e., the alert profile schema may be a collection of entities with types) and may leverage an entity-relation data model. As will be described in more detail below, alert profiles may help policies define certain aspects of how alerts are to be provided, displayed, prioritized, highlighted, and so forth so as to ensure that user 517 of user interface 516 receives, and is able to successfully take notice of and understand, each of various different alerts that may be generated as part of incoming query results such as query result 514.

Just as queries and policies described above may provide various benefits by being modular and customizable (e.g., such that different policies may invoke the same query, the same policy may invoke various different queries, policies and queries may both be provided by the data platform by default or may be customized to the needs of a client entity, etc.), so too may alert profiles be implemented in flexible, customizable, modular, and extensible ways. For example, an alert profile authored by a particular person associated with a particular client entity may be modular so as to be deployed with (i.e., invoked by) a policy authored by a different person associated with the particular client entity or even by a person associated with a different client entity. Moreover, alert profiles may be extensible in the sense that each invocation of a particular alert profile may have its own configuration override (i.e., the attributes of a given alert profile may be overridden at each policy reference) and the particular alert profile may be a subclass of another alert profile by leveraging an inheritance mechanism supported by the query language in which the alert profile is developed.

These and various other aspects and features related to alert profiles will now be described in relation to FIGS. 6A-6E. In particular, an ability of the data platform to associate the same alert profile with multiple custom policies created using the custom query language for the data platform will be described. Various benefits of this approach to alert profiles will be made apparent in the following description, including that users may draft policies and/or queries without necessarily having to understand and consider reporting and alerting details handled by modular alert profiles that may have been drafted and documented by other parties and that are available for use with new custom policies.

FIG. 6A illustrates an example of a process 600 for implementing modular alert profiles for custom policies and queries. While FIG. 6A shows illustrative operations according to one implementation, other implementations may omit, add to, reorder, and/or modify any of the operations shown in FIG. 6A. In some examples, multiple operations shown in FIG. 6A or described in relation to FIG. 6A may be performed concurrently (e.g., in parallel) with one another, rather than being performed sequentially as illustrated and/or described. One or more of the operations shown in FIG. 6A may be performed by a data platform described herein (e.g., data platform 12). More particularly, computing resources of a data platform (e.g., data ingestion resources 18, data processing resources 20, etc.) may perform each of the operations of process 600 during or after a data ingestion procedure has been performed to move a target dataset from a cloud environment (e.g., cloud environment 14) to a data store (e.g., data store 30). The operations of process 600 may be performed prior to, in parallel with, or subsequent to other operations described above to be performed by the data platform. For example, operations 601-604 may be performed before, during, or after the data platform constructs a graph (e.g., a polygraph) in any of the ways described herein. Based on client application data ingested from the cloud environment (e.g., cloud environment 14) and loaded into the data store (e.g., data store 30), the data platform (e.g., data platform 12) may construct a graph (e.g., a polygraph) comprising a plurality of nodes connected by a plurality of edges. As described above, each node of the plurality of nodes may represent a logical entity from the client application data and each edge of the plurality of edges may represent a behavioral relationship between nodes connected by the edge.

Each of operations 601-604 of process 600 will now be described in more detail as the operations may be performed by computing resources of a data platform described herein. The graph construction and/or other such operations described above may be performed in connection with operations 601-604 in any suitable way.

At operation 601, the data platform (e.g., computing resources operating on the data platform) may determine that a particular alert profile is invoked by a first custom policy. For example, the first custom policy may reference or link to the particular alert profile so as to rely on code or other aspects of the alert profile to handle alerting and reporting details and other contract constraints with services such as an event generator service. The first custom policy may further invoke a first query associated with the data platform. For example, the particular alert profile and the first query may be invoked in the ways described and illustrated above, and, as with various queries that have been described, the first query may target a target dataset ingested from a cloud environment and stored in a data store.

At operation 602, the data platform may, similarly to operation 601, determine that the particular alert profile (i.e., the same particular alert profile invoked by the first custom policy at operation 601) is also invoked by a second custom policy that is distinct from the first custom policy. Similar to the first custom policy, the second custom policy may further invoke a second query associated with the data platform, where the second query, like the first query and other queries described herein, may target the target dataset stored in the data store. While the second custom policy invoking the particular alert profile is distinct from the first custom policy, it will be understood that the particular alert profile being invoked is a same modular alert profile in both cases. It will also be understood that the first and second queries invoked by the respective first and second custom policies may be separate and distinct queries in certain examples, or may be one and the same query in other examples.

At operation 603, the data platform may annotate a first query result generated based on the first query. Specifically, the annotating of the first query result at operation 603 may be performed in accordance with the particular alert profile due to (or as a result of) the invoking of the particular alert profile in the first custom policy. Examples of alert profiles used to define the annotating of query results will be described and illustrated in more detail below.

Similarly, at operation 604, the data platform may annotate a second query result generated based on the second query. Again, the annotating of the second query result may also be performed in accordance with the particular alert profile, in this case due to (or as a result of) the invoking of the particular alert profile in the second custom policy.

FIG. 6B illustrates an example of how different custom policies may invoke a same alert profile. Specifically, in FIG. 6B two different client entities 610 (i.e., a first client entity 610-1 and a second client entity 610-2) each represented by a particular user 611 (i.e., a first user 611-1 associated with client entity 610-1 and a second user 611-2 associated with client entity 610-2). Client entities 610 are shown on either side of a dashed line representing authorship of various code. Specifically, in this example, user 611-1 of client entity 610-1 will be understood to have authored certain custom policies 612 that are above the dashed line (i.e., custom policies 612-1 and 612-2) as well as a query 613 that those policies each invoke (i.e., query 613-1). Meanwhile, user 611-2 of client entity 610-2 will be understood to have authored other custom policies 612 below the dashed line (i.e., custom policies 612-3, 612-4,

and **612-5**) as well as queries **613** that these policies respectively invoke (i.e., queries **613-2**, **613-3**, and **613-4**). Similar to custom policies and queries described above, a respective query result **614** having certain annotations **615** is shown to be generated in accordance with each of custom policies **612**. Specifically, as shown, a query result **614-1** with annotations **615-1** is provided in accordance with custom policy **612-1**, a query result **614-2** with annotations **615-2** is provided in accordance with custom policy **612-2**, a query result **614-3** with annotations **615-3** is provided in accordance with custom policy **612-3**, and so forth for query results **614-4** and **614-5** and respective annotations **615-4** and **615-5**.

An alert profile **616-1** is shown above the dashed line to have been authored by user **611-1** of client entity **610-1**, and this alert profile is invoked by several custom policies **612** including custom policy **612-3**, which was authored by the other client entity (i.e., client entity **610-2**). Alert profile **616-1** is also shown to inherit or extend an antecedent alert profile **617** (illustrated above the dashed line but understood to have been authored by any client entity). Another alert profile **616-2** is shown below the dashed line to be invoked by another custom policy (custom policy **612-4**), while yet another custom policy shown in FIG. 6B (custom policy **612-5**) is shown to not invoke any alert profile at all (illustrated as a reference to a null alert profile **618**). Each of the aspects of FIG. 6B will now be described in more detail.

Client entities **610** may refer to any suitable client entities as may serve a particular implementation. As mentioned above, a client entity for a data platform, as those terms are used herein, may refer to a client or customer company or other organization that uses (e.g., receives services from) the data platform. As such, client entities may be understood to be organizations who manage cloud data deployed to cloud environments described herein using services (e.g., security services, graphing services, querying services, etc.) provided by the data platform. Each client entity **610** may be associated with (e.g., may employ, etc.) various users such as data security professionals, information technology experts, and so forth. Accordingly, users **611** may represent one particular user associated with each of the illustrated client entities **610**, though users **611** are understood to serve as illustrative representations of what may actually be entire teams working for a given client entity. In the course of their work, users **611** may draft the various custom policies, queries, alert profiles, and other code relevant to services of the data platform. In some examples, as will be described in more detail below, that code may be drafted to be portable and modular (e.g., usable by others), extensible (e.g., readily built upon and extended), reusable, flexible, and to possess other desirable qualities.

Custom policies **612** may represent any of the custom policies (or in certain examples, default policies) described herein. As such, custom policies **612** may be implemented in accordance with any of the examples provided above (see, e.g., FIGS. 5E-5K) and may invoke queries **613** and/or alert profiles **616** by referencing them in the ways that have been shown and described.

Queries **613** too may be implemented in accordance with example queries that have been described, and query results **614** and annotations **615** may be generated and provided accordingly. As shown, in certain examples, a first query invoked by a first custom policy (e.g., custom policy **612-1**) and a second query invoked by second custom policy (e.g., custom policy **612-2**) may be a same query (e.g., query **613-1** in this example). In other examples, however, it may also be the case that a first query **613** invoked by a first

custom policy (e.g., custom policy **612-3**, which invokes query **613-2**) may be different from a second query **613** invoked by a second custom policy (e.g., custom policy **612-4**, which invokes query **613-3**).

Query results **614** and annotations **615** may be implemented in the same way as other query results and annotations described herein. However, it is emphasized for purposes of FIG. 6B that different query results **614** may be annotated with respective annotations **615** in accordance not only with code incorporated into the respective custom policy **612**, but also in accordance with the particular alert profile **616** that the custom policy invokes. For example, since both custom policies **612-1** and **612-2** invoke alert profile **616-1**, annotations **615-1** and **615-2** or query results **614-1** and **614-2** (respectively) may each reflect various aspects defined in alert profile **616-1** with respect to how alerts are to be handled, presented, prioritized, and so forth. Indeed, whenever a custom policy **612** is configured to generate alerts in addition to (or as part of) the query results that it outputs, a reference to a particular alert profile may be required so as to set forth various parameters for how the alerts are to be generated and provided.

Multiple custom policies may reference the same alert profile **616** (e.g., illustrated in FIG. 6B by custom policies **612-1**, **612-2**, and **612-3** each invoking a same alert profile **616-1**). As shown, this may be the case even if a modular alert profile happens to have been authored by a different user **611** or client entity **610** than the one that authored the custom policy. More specifically, a particular alert profile (e.g., alert profile **616-1**) may be developed by a first client entity (e.g., client entity **610-1**) that is distinct from a second client entity (e.g., client entity **610-2**) that develops at least one of the custom policies **612** that invoke that particular alert profile (e.g., custom policy **612-3**).

Along with the modularity and reusability of alert profile **616-1** illustrated by its invocation by several different custom policies **612**, alert profiles such as alert profile **616-1** may be extensible in the sense that they may benefit by inheritance from existing alert profiles such as antecedent alert profile **617**. Specifically, a particular alert profile such as alert profile **616-1** may reference an antecedent alert profile (e.g., antecedent alert profile **617**), which the particular alert profile may be configured to extend by both 1) inheriting a first attribute (or first set of several attributes) included in the antecedent alert profile, and 2) contributing a second attribute (or second set of several attributes) that are absent from the antecedent alert profile (i.e., attributes that are new, rather than being inherited). Examples of how an alert profile may reference and extend an antecedent alert profile will be described in more detail below.

Another concept illustrated by FIG. 6B is that not all policies require alert profiles since certain policies do not even result in alerts. As shown, along with various custom policies **612-1** through **612-4** that are understood to include respective alert attributes that are enabled (thereby causing each of these to invoke a non-null alert profile **616**), the data platform may also determine that no alert profile is invoked by another custom policy such as custom policy **612-5**. More specifically, as shown, whereas the other custom policies **612** invoke alert profiles **616**, custom policy **612-5** is shown to invoke null alert profile **618**, which will be understood to be equivalent to not invoking any alert profile at all (as indicated by the dotted line outlining null alert profile **618**). Whereas custom policies **612-1** through **612-4** were each described as having respective alert attributes that were enabled, custom policy **612-5** will be understood to include its own alert attribute that is disabled. For example, referring

to the custom policy examples described above in relation to FIGS. 5E-5K, the “alertEnabled” setting of custom policy 612-5 may be set to “false” (rather than “true”) and the “alertProfile” setting of custom policy 612-5 may be set to “null” rather than referencing an actual alert profile.

FIG. 6C illustrates an example architecture 620 by way of which custom policies invoking a same alert profile may generate alerts for presentation to a user. As shown, various custom policies 612 (any of custom policies 612 described above as being authored by any user or client entity) may invoke various queries 613 (also authored by arbitrary users and client entities and potentially being the same or different queries invoked by various custom policies 612, as described above). Additionally, a same alert profile 616 is shown in FIG. 6C to be invoked by the various custom policies 612. Along with these familiar elements, architecture 620 also introduces a policy evaluator 621 configured to generate evaluation context data 622 (e.g., including observations, etc., regarding alert events arising based on custom policies 612, queries 613, and alert profile 616); an event generation service 623 configured to generate a set of alerts 624 (labeled as alerts 624-1 at this stage of the architecture); and an event triage service 625 configured to output the set of alerts 624 (labeled as alerts 624-2 at this stage of the architecture) to the user interface.

Policy evaluator 621 may be configured to evaluate query results of queries 613 based on the custom policies 612 and the alert profile 616. For example, policy evaluator 621 may verify that an alert attribute of each policy is enabled (e.g., alertEnabled is set to true) and that alert profile 616 is valid (e.g., alertProfile is set to reference an actual alert profile rather than “null”). Assuming everything is properly enabled and referenced, policy evaluator 621 will fetch the invoked alert profile (alert profile 616) and determine if default values are to be used or if any override has been defined (overrides will be described in more detail below). Policy evaluator 621 then generates, based on the evaluating of the query results, evaluation context data 622 that may include an observation with respect to at least one of the query results associated with one of custom policies 612. This data (including the observations) will be used as an input by the event generation service to create alerts that will eventually be presented to the user.

Event generation service 623 may be configured to generate, based on evaluation context data 622, a plurality of alerts 624-1 that each indicate a respective observation included within evaluation context data 622. Event generation service 623 may implement EventGen 158 described above. As such, event generation service 623 may examine observations in aggregate to combine observations of a similar type into a single alert, to deduplicate observations (and, in certain examples, increase the severity of the remaining observation as a result of dropping the duplicate), to apply suppression rules prior to generating final alerts, to score alerts being generated (e.g., readjusting the severity based on custom rules or scoring algorithms), to perform data enrichment or threat enrichment, to compute data that may assist a user reviewing the alert, and so forth. Various examples of events that may be detected and alerted on by an event generation service are described above in relation to EventGen 158. After generating the alerts, event generation service 623 may then provide these alerts 624-1 for presentation to a user of the data platform.

In certain examples, as shown, before being provided to the user interface for presentation to a user, alerts 624 may first be processed by an event triage service 625. Event triage service 625 may be configured to prioritize, based on

the respective observations indicated by the plurality of alerts, certain alerts over one or more other alerts of the plurality of alerts 624-1 received from event generation service 623. As such, the providing of any particular alert may be performed based on the prioritizing of the alert over the one or more other alerts at event triage service 625. Accordingly, as shown, alerts 624-1 that are generated by event generation service 623 may be extracted, enriched, prioritized, and otherwise processed by event triage service 625 such that alerts 624-2 may be provided to the user interface for presentation to the user (e.g., by way of the API, as has been described).

As mentioned above (and as illustrated by antecedent alert profile 617 and alert profile 616-1 in FIG. 6B), a particular alert profile may be configured to reference an antecedent alert profile that the particular alert profile is configured to extend by both 1) inheriting a first attribute that is included in the antecedent alert profile; and 2) contributing a second attribute that is absent from the antecedent alert profile. To illustrate, FIG. 6D shows alert profiles 631, 632, and 633 to illustrate example alert profiles and demonstrate how inheritance between alert profiles may be implemented. In FIG. 6D and other figures herein, sample code may be written using a monospace font that is intermixed with a proportional font that indicates (within comments indicated between escape characters ‘/*’ and ‘*/’) non-code descriptions of code that may be present in a given example, but which is omitted from the figure due to space constraints. For example, as shown in FIG. 6D, alert profile 631 is titled “LW_Profile” and includes code defining certain “fields” and “entityKeys.” However, rather than indicating the precise code for the various fields and entity keys that may be included in this example, the proportional font indicates that, between the brackets following the “fields:”, the code may “define fields such as Number, String, Timestamp, IP_ADDR, PORT, PROTOCOL, AWS_ACCOUNT_ID, etc.” Similarly, between the brackets following the “entityKeys:”, the code (not explicitly shown) will be understood to “define entity keys such as ProcessKey, MachineKey, ConnectionKey, etc.”

Alert profile 631 may be understood to serve as an antecedent alert profile to alert profile 632, is indicated by the code in alert profile 632 stating that this alert profile (titled “LW_AlertProfile”) “extends: LW_Profile” (i.e., is an extension of alert profile 631). Alert profile 632 is shown to also define certain fields, some of which may be inherited from LW_Profile (i.e., alert profile 631) and others of which may be newly defined and not inherited. Though alert profile 632 is itself extended from an antecedent alert profile (alert profile 631), alert profile 632 may also be extended further by other alert profiles that use it as an antecedent alert profile. To illustrate, an alert profile 633 titled “LW_AWS_AlertProfile” includes code to indicate that it “extends: LW_AlertProfile”, such that attributes of both LW_AlertProfile, as well as LW_Profile, may all automatically be incorporated into alert profile 633 along with new attributes that alert profile 633 defines.

As shown, alert profile 633 may extend and define new fields and entity keys.

Additionally, alert profile 633 is shown to define certain “descriptionKeys” and to define at least one alert. As shown, the alert defined by alert profile 633 includes a name (“NewViolation”), an event type (“New Violations”), an event name (“NewViolations”), and a description and subject that each may reference various description keys that have been defined and/or other fields, keys, or other suitable data as may serve a particular implementation. While only a

single alert is explicitly shown to be defined in alert profile **633**, it will be understood that more than one alert may be defined and the first may be used as a default when alert profile **633** is invoked.

The inheritance illustrated in FIG. **6D** facilitates modularity, portability, and extensibility for alert profiles that may be authored for one custom policy or query and subsequently used for another. Another similar feature that further facilitates these desirable features is referred to herein as an override. When a particular alert profile is nearly, but not quite, optimal for a given use case, a user may use the particular alert profile as an antecedent alert profile that forms a basis for a new alert profile (which inherits the desirable attributes as described above), or the user may invoke the non-optimal alert profile directly using an override to adjust the undesirable or non-optimal attributes.

To illustrate, FIG. **6E** shows example custom policies **634** and **635** that invoke a particular alert profile with and without overriding an attribute of the alert profile. Specifically, as shown, custom policy **634** invokes “LW_AWS_AlertProfile” (i.e., alert profile **633** of FIG. **6D**) as-is (without an override), while custom policy **635** invokes this same LW_AWS_AlertProfile with an override **636** that adjusts certain attributes to better serve the objectives of custom policy **635**.

Given that different custom policies may invoke the same alert profile, as has been described (see, e.g., FIGS. **6A-6B**), the policies of FIG. **6E** (without and with override **636**) illustrate at least three potential scenarios for a first custom policy and a second custom policy that each invoke the same alert profile. Each of these three scenarios will now be described.

First, both the first and the second custom policies may invoke the particular alert profile as-is, without overriding any attribute of the particular alert profile. For example, in this scenario, the first custom policy may be implemented by custom policy **634** and the second custom policy may be implemented by a different custom policy that likewise invokes LW_AWS_AlertProfile without any override.

Second, the first custom policy may invoke the particular alert profile as-is (i.e., without overriding any attribute of the particular alert profile), while the second custom policy may invoke the particular alert profile with an override configured to change one or more attributes of the particular alert profile (for more optimal annotating of the second query result that will be annotated based on the second custom policy). For example, in this scenario, the first custom policy may be implemented by custom policy **634** (as-is, with no override) and the second custom policy may be implemented by custom policy **635** (including override **636**, which changes certain attributes such as the fields and field-`TypeMap`).

Third, the first custom policy may invoke the particular alert profile with a first override configured to change an attribute of the particular alert profile in a first way (e.g., for more optimal annotating of the first query result that will be annotated based on the first custom policy), while the second custom policy may invoke the particular alert profile with a second override configured to change the attribute of the particular alert profile in a second way that is different from the first way (e.g., for more optimal annotating of the second query result that will be annotated based on the second custom policy). For example, in this scenario, the first custom policy may be implemented by custom policy **635** including override **636**, while the second custom policy may be implemented by another custom policy similar to custom policy **635** but with an override that is distinct from override

636 (e.g., by changing different attributes, by changing the same attributes in a different way, a combination of these, etc.).

FIGS. **6A-6E** have illustrated various scenarios and aspects related to modular alert profiles that may be used not only by a user or client entity that originally authors and fully understands the alert profiles, but also potentially by other users (including users associated with other client entities) who are authoring their own custom policies and may not wish to draft or come to fully understand various details regarding how the modular profiles work. Significant advantages and benefits have been described and made apparent for alert profiles that can be modularly shared and used according to this type of paradigm.

With reference to FIGS. **7A-7C**, a somewhat similar type of sharing and modularity will now be described, this one relating to custom policies themselves. Specifically, FIGS. **7A-7C** illustrate various aspects of community-based generation of policies for the data platform (e.g., community-driven creation of custom policies). Client entities relying on data platforms described herein may be very sophisticated with regard to their data, including how and where the data is stored, managed, secured, kept in compliance, monitored, and so forth. In certain cases, a client entity may employ large teams of professionals dedicated to securing and strategically managing the valuable data owned by the client entity, and these teams may be well-informed and highly knowledgeable about how to intelligently and efficiently handle various scenarios (e.g., security threats, various types of non-compliance, different types of data events, etc.). Accordingly, not only is it desirable to provide such professionals with the flexible data management tools that have been described (e.g., custom policies, custom queries, modular alert profiles, etc.), but it may also benefit both the data platform and the various client entities using the data platform to leverage the various insights, knowledge, creativity, and experience from the community of client entities making use of these tools in intelligent and creative, yet potentially distinct and innovative, ways.

Methods and systems for community-based generation of policies will now be described to allow data platforms to implement a “virtuous cycle” of providing effective policies to competent and sophisticated client entities (along with flexibility to customize those policies and/or create new ones), then learning from the customized policies of those client entities to improve and expand the policy offering. In this way, the data platform may grow and improve to better serve all of the client entities. While community-based generation of policies will be the focus of the following description, it will be understood that this same principles may also be applied to other aspects of the customer experience. For example, community-based generation of queries (e.g., custom queries), modular alerts, user interface views, event generation and triage parameters, and so forth, may each be implemented to similarly beneficial results.

As will be described in more detail below, community-based generation of policies may be performed and/or implemented in any manner as may serve a particular implementation. For example, the data platform may keep track of common custom policies and/or queries that various client entities use and may compile and maintain a list of these to be offered for use by the various client entities (e.g., all client entities using the data platform, certain client entities that meet specific criteria, etc.). Policies and/or queries on such a list may be automatically or manually (e.g., under supervision of a person responsible for managing such data) brought into a repository of what are referred

to herein as “community policies” or “community queries.” In some implementations, such community policies and queries may be functionally equivalent to default policies and queries (in that they are provided to some or all client entities by the data platform itself), while in other implementations there may be significant differences such that community policies and queries are maintained and offered separately from default policies and queries.

In certain examples, custom policies and queries that are deemed to be of general usefulness may be designated as-is (i.e., in their original form) to be used as community policies, and to be offered for use by users associated with client entities other than the client entity that originally developed (authored) the policies and queries. In other examples, popular or useful custom policies and queries may not necessarily be designated and offered up as community policies as-is, but rather may form a basis for community policies that are generated by the data platform. For example, a genericized version of a popular custom policy (e.g., a version that removes or generalizes client-specific parameters from the original, etc.) may be created and offered up as a community policy; certain aspects of a plurality of similar custom policies may be combined to form a single community policy; a particular feature of a custom query (e.g., the filter defining the aspect that is to be identified) may be used for a community query that targets a different dataset or exposes different fields in its query results; one or more common custom queries may be used to create a flexible template including certain predetermined parameters and others parameters open to be customized by new users; and so forth.

FIG. 7A illustrates an example of a process 700 for community-based generation of policies for a data platform. While FIG. 7A shows illustrative operations according to one implementation, other implementations may omit, add to, reorder, and/or modify any of the operations shown in FIG. 7A. In some examples, multiple operations shown in FIG. 7A or described in relation to FIG. 7A may be performed concurrently (e.g., in parallel) with one another, rather than being performed sequentially as illustrated and/or described. One or more of the operations shown in FIG. 7A may be performed by a data platform described herein (e.g., data platform 12). More particularly, computing resources of a data platform (e.g., data ingestion resources 18, data processing resources 20, etc.) may perform each of the operations of process 700 during or after a data ingestion procedure has been performed to move a target dataset from a cloud environment (e.g., cloud environment 14) to a data store (e.g., data store 30). The operations of process 700 may be performed prior to, in parallel with, or subsequent to other operations described above to be performed by the data platform. For example, operations 701-703 may be performed before, during, or after the data platform constructs a graph (e.g., a polygraph) in any of the ways described herein. Based on client application data ingested from the cloud environment (e.g., cloud environment 14) and loaded into the data store (e.g., data store 30), the data platform (e.g., data platform 12) may construct a graph (e.g., a polygraph) comprising a plurality of nodes connected by a plurality of edges. As described above, each node of the plurality of nodes may represent a logical entity from the client application data and each edge of the plurality of edges may represent a behavioral relationship between nodes connected by the edge.

Each of operations 701-703 of process 700 will now be described in more detail as the operations may be performed by computing resources of a data platform described herein.

The graph construction and/or other such operations described above may be performed in connection with operations 701-703 in any suitable way.

At operation 701, the data platform (e.g., computing resources operating on the data platform) may access a custom policy developed by a particular client entity that uses the data platform. As with various custom policies described herein, the custom policy accessed at operation 701 may invoke a query (e.g., a default or custom query) that targets a target dataset ingested from a cloud environment and stored in a data store. For example, as has been described, the target dataset may include trail data (e.g., AWS CloudTrail data or the like), client application data, and/or any other data as may have been ingested from the cloud environment to the data store in a particular implementation.

At operation 702, the data platform may determine that the custom policy is likely to be of value to one or more other client entities, besides the particular client entity, that also use the data platform. For example, the data platform may determine that the custom policy is designated for publishing (or as being of general usefulness) by its author, that the custom policy is used frequently, that the custom policy is applied to different sets of circumstances in a general way, that the custom policy is highly ranked or produces results that tend to be prioritized and/or highly scored, or that the custom policy otherwise exhibits any suitable factor tending to indicate that the policy is of general usefulness or exceptional value.

At operation 703, the data platform may generate a community policy based on the custom policy. The generating of the community policy at operation 703 may be performed in response to the determining at operation 702 that the custom policy is likely to be of value to the one or more other client entities. As mentioned above, in certain examples, the generating of the community policy may include designating the custom policy as a community policy as-is (e.g., without necessarily changing the custom policy from its original form). In other examples, the community policy may be generated based on the custom policy in a manner that results in at least some changes to the custom policy (e.g., the removal of client-specific parameters or data, the generalizing of certain terms, etc.). In any case, the community policy generated at operation 703 may be made available for use by the one or more other client entities besides the particular client entity that originally developed the custom policy.

FIG. 7B illustrates an example configuration in which community-based generation of custom policies is implemented. Specifically, as shown, several different users 711 (e.g., users 711-1 through 711-3) associated with respective client entities 712 (e.g., client entities 712-1 through 712-3) are shown to have authored respective custom policies 713 (e.g., custom policies 713-1 through 713-3). More particularly, as illustrated by dotted lines extending between the client entities 712 and the custom policies 713, user 711-1 of client entity 712-1 will be understood to have created custom policy 713-1, user 711-2 of client entity 712-2 will be understood to have created custom policy 713-2, and user 711-3 of client entity 712-3 will be understood to have created custom policy 713-3. Users 711, client entities 712, and custom policies 713 may be examples of (or may be implemented by) any of the users, client entities, and/or custom policies that have been described herein.

Each custom policy 713 is shown to invoke a query 714, and these invocations will be understood to be performed in accordance with any of the examples provided above. As

with other custom policies described herein, each custom policy 713 may uniquely invoke a particular query not invoked by other policies or may be one of a plurality of policies invoking a same query. Moreover, each query invoked by one of custom policies 713 may be either a custom query (e.g., a custom query that is also developed by the same client entity 712 that authored the custom policy, a custom query developed by a different client entity 712, etc.) or a default query that is provided by the data platform to a plurality of client entities (e.g., to all the client entities 712 including the particular client entity that authored the custom policy). To illustrate some of this possible variation, the configuration of FIG. 7B shows that both custom policies 713-1 and 713-2 invoke a same custom query 714-1. Custom query 714-1 will be understood to have been developed by client entity 712-1 (e.g., by user 711-1), such that it is developed by the same client entity as developed custom policy 713-1 but by a different client entity as developed custom policy 713-2. The configuration of FIG. 7B also shows that custom policy 713-3 alone invokes a default query 714-2 that is provided by the data platform.

FIG. 7B also illustrates a policy manager 715 implemented by computing resources of the data platform and within which a machine learning model 716 is deployed. Arrows indicate that any or all of custom policies 713 may be considered by policy manager 715 (e.g., using machine learning model 716) to determine if the custom policies are likely to be of value to client entities besides the client entities that originally developed them (i.e., operation 702 of process 700 in FIG. 7A). While no such arrow is shown in this configuration to indicate that custom query 714-1 may similarly be considered as a basis for community-based content, it will be understood that, in certain implementations, custom queries such as custom query 714-1 (and other custom queries described herein) may undergo a process analogous to what is being illustrated and described here for custom policies.

Policy manager 715 may access custom policies 713 and determine whether they are likely to be of value generally (e.g., to client entities besides the client entities that authored or developed the custom policies) in any manner as may serve a particular implementation. For example, certain implementations of policy manager 715 may access and analyze all custom policies 713 that are implemented using the data platform and give them due consideration for use in generating community policies. Other implementations of policy manager 715 may access and analyze only custom policies 713 that are designated (e.g., by their authors) for such consideration (e.g., by being voluntarily published or shared).

The determination of whether a given custom policy is likely to be of value to other client entities (and therefore may be designated as, or used as a basis for creating, a community policy) may also be performed in any suitable way. For instance, the determining that the custom policy is likely to be of value to the one or more other client entities may be performed based on an identity of the particular client entity (e.g., whether the authoring client entity is known to be particularly sophisticated, whether the authoring client entity's data management practices are considered to be generalizable to other client entities, etc.); based on a popularity metric of the custom policy exceeding a predetermined popularity threshold (e.g., in an implementation in which popularity metrics for different custom policies are tracked and maintained based on usage and/or user feedback); based on user input representative of a request to publish the custom policy as a community policy (e.g., in

implementations in which the policy manager receives such user input representative of the request to publish the custom policy from the particular client entity); or based on other suitable conditions, circumstances, or events.

To illustrate, FIG. 7C illustrates example aspects of how policy manager 715 may make the determination that a custom policy 713 is likely to be of value to one or more other client entities and, more generally, how a community policy may be generated based on the custom policy 713. As shown in FIG. 7C, a particular custom policy 713 may include metadata indicating, for example, an author 721 of the custom policy, a popularity metric 722 of the custom policy, a publish flag 723 indicating whether the custom policy is requested or approved for publishing as a community policy, and other metadata as may serve a particular implementation.

In some implementations, custom policies may be published and shared outside of a particular user and/or client entity without necessarily being considered for or used in the generation of a community policy. For instance, a user may share a particular custom policy with a coworker associated with the same client entity or with a colleague at another client entity. In these examples, the data platform (and, specifically, policy manager 715 operating within the data platform) may track a popularity metric for this shared custom policy to determine whether the shared custom policy should be fully adopted for use in generating one or more community policies. For example, the popularity metric in this case could be determined based on how many other users (or client entities) have used a particular shared policy, how many instances of the shared policy other users (or client entities) have deployed, or the like. Accordingly, a certain popularity threshold (e.g., a threshold number of other users who have deployed a policy, a threshold number of instances in which the policy has been deployed, etc.) may be compared against the popularity metric to help make the determination of whether the shared policy is likely to be of value if offered more widely as a community policy.

Upon receiving custom policy 713 and accompanying metadata such as author 721, popularity metric 722, and publish flag 723, policy manager 715 may generate a community policy 720 based on custom policy 713, and may provide this community policy 720 to a policy repository 717 (described in more detail below). As mentioned above, community policy 720 may be similar or identical to custom policy 713 (e.g., such as if custom policy 713 is designated as-is to serve as a community policy), or may be distinct from, but still based on, custom policy 713 (e.g., such as if custom policy 713 is significantly changed to genericize its parameters or if aspects of custom policy 713 are combined with one or more other custom policies to generate a hybrid policy for community policy 720).

Returning to FIG. 7B, policy manager 715 is also shown to include machine learning model 716, which may be used in addition, or as an alternative, to the determination methods described above. Specifically, in implementations that make use of such a model, the determining that the custom policy is likely to be of value to the one or more other client entities may be performed automatically using machine learning model 716. Machine learning model 716 may be trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities. In this way, with or without the metadata described in relation to FIG. 7C, policy manager 715 may become progressively adept (as machine learning model 716

continues to learn and develop) at determining which custom policies are likely to make useful community policies and which are less likely to do so.

As mentioned above, community policies generated by policy manager 715 (e.g., community policy 720, etc.) may be maintained in a policy repository 717. More specifically, along with analyzing custom policies and generating community policies, policy manager 715 (or another suitable component of the data platform) may manage policy repository 717 as a storage repository configured to store a plurality of community policies that are available for use by a plurality of different client entities, and may store (e.g., in response to the generating of the community policy) each generated community policy (e.g., including community policy 720) within policy repository 717. Dashed lines connecting each client entity 712 to policy repository 717 are drawn to represent the access that various users 711 of the client entities 712 may have to the policies of policy repository 717 (i.e., the availability of those policies for use by those users). In this example, as shown, all of users 711 and client entities 712 may access any of the policies managed in policy repository 717. In other examples, certain policies may be restricted to only certain users and/or client entities as may serve a particular implementation.

In the example of FIG. 7B, different sections of policy repository 717 are dedicated 1) to default policies that are developed and provided by the data platform and 2) to community policies generated by policy manager 715 based on custom policies drafted by clients as described above. It will be understood that this illustrated structure is illustrative only, and that certain implementations of a policy repository may include more or fewer types of data instances than shown here. For instance, certain policy repositories may be dedicated only to default policies or only to community policies (rather than combining both). As another example, certain repositories may include not only default and/or community policies, but also default and/or community queries, default and/or community alert profiles, and/or other suitable data instances that may be accessible to multiple client entities.

As shown by “Documentation” illustrated alongside each default and community policy in policy repository 717, the managing of policy repository 717 may include storing documentation corresponding to each community and/or default policy stored within policy repository 717, as well as providing documentation corresponding to a particular community or default policy to a particular client device (e.g., device used by one of users 711) accessing the particular community or default policy from policy repository 717. The documentation provided with the policies in this way may lay out any suitable features, parameters, or other information about the associated policy as may be useful to a user reading the documentation. As such, an understanding of basic concepts from the documentation may allow a user to leverage policies even if he or she has not studied or fully grasped the code of the policy itself.

What is claimed is:

1. A method comprising:

accessing, by a data platform, a custom policy developed by a particular client entity that uses the data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store;

determining, by the data platform, that the custom policy is likely to be of value to one or more other client entities based on a combination of one or more of:

a request to publish the policy,

a popularity threshold based on usage, a sophistication of the author, and user feedback; and

generating, by the data platform and in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities, wherein the custom policy is a compliance policy configured to verify that the target dataset complies with a predetermined compliance protocol.

2. The method of claim 1, wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on an identity of the particular client entity.

3. The method of claim 1, wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on a popularity metric of the custom policy exceeding a predetermined popularity threshold.

4. The method of claim 1, further comprising:

receiving, by the data platform from the particular client entity, user input representative of a request to publish the custom policy as a community policy;

wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on the user input representative of the request to publish the custom policy.

5. The method of claim 1, wherein:

the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and

the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.

6. The method of claim 1, further comprising:

managing, by the data platform, a policy repository configured to store a plurality of community policies that are available for use by a plurality of different client entities; and

storing, by the data platform and in response to the generating of the community policy, the community policy within the policy repository.

7. The method of claim 6, wherein the managing of the policy repository includes:

storing documentation corresponding to each community policy stored within the policy repository; and

providing documentation corresponding to a particular community policy to a particular client device accessing the particular community policy from the policy repository.

8. The method of claim 1, further comprising providing, by the data platform and for use by a user of the data platform, a graphical user interface that provides graphical assistance configured to assist a user in performing one or more of:

discovering default policies and community policies available for use by a client entity with which the user is associated;

creating new custom policies; or

requesting to publish custom policies as community policies that are available for use by other users associated

101

with other client entities besides the client entity with which the user is associated.

9. The method of claim 1, wherein the query invoked by the custom policy is a custom query that is also developed by the particular client entity.

10. The method of claim 1, wherein the query invoked by the custom policy is a default query provided by the data platform to a plurality of client entities including the particular client entity.

11. The method of claim 1, further comprising constructing, by the data platform and based on client application data ingested from the cloud environment and loaded into the data store, a graph comprising a plurality of nodes connected by a plurality of edges, wherein each node of the plurality of nodes represents a logical entity from the client application data and each edge of the plurality of edges represents a behavioral relationship between nodes connected by the edge.

12. A system comprising:

a memory storing instructions; and
a processor communicatively coupled to the memory and configured to execute the instructions to:

access a custom policy developed by a particular client entity that uses a data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store;

determine that the custom policy is likely to be of value to one or more other client entities based on a combination of one or more of:

a request to publish the policy,
a popularity threshold based on usage,
a sophistication of the author, and
user feedback; and

generate, in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities, wherein the custom policy is a compliance policy configured to verify that the target dataset complies with a predetermined compliance protocol.

13. The system of claim 12, wherein:

the processor is further configured to execute the instructions to receive, from the particular client entity, user input representative of a request to publish the custom policy as a community policy; and

the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on at least one of:

an identity of the particular client entity,
a popularity metric of the custom policy exceeding a predetermined popularity threshold, or
the user input representative of the request to publish the custom policy.

14. The system of claim 12, wherein:

the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and

the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.

102

15. The system of claim 13, wherein the processor is further configured to execute the instructions to:

manage a policy repository configured to store a plurality of community policies that are available for use by a plurality of different client entities; and

store, in response to the generating of the community policy, the community policy within the policy repository.

16. A computer program product embodied in a non-transitory computer readable storage medium and comprising computer instructions for:

accessing a custom policy developed by a particular client entity that uses a data platform, wherein the custom policy invokes a query that targets a target dataset ingested from a cloud environment and stored in a data store;

determining that the custom policy is likely to be of value to one or more other client entities based on a combination of one or more of:

a request to publish the policy,
a popularity threshold based on usage,
a sophistication of the author, and
user feedback; and

generating, in response to the determining that the custom policy is likely to be of value to the one or more other client entities, a community policy based on the custom policy, wherein the community policy is available for use by the one or more other client entities, wherein the custom policy is a compliance policy configured to verify that the target dataset complies with a predetermined compliance protocol.

17. The computer program product of claim 16, further comprising computer instructions for receiving, from the particular client entity, user input representative of a request to publish the custom policy as a community policy;

wherein the determining that the custom policy is likely to be of value to the one or more other client entities is performed based on at least one of:

an identity of the particular client entity,
a popularity metric of the custom policy exceeding a predetermined popularity threshold, or
the user input representative of the request to publish the custom policy.

18. The computer program product of claim 16, wherein: the determining that the custom policy is likely to be of value to the one or more other client entities is performed automatically using a machine learning model implemented by the data platform; and

the machine learning model is trained for identifying custom policies of value to multiple client entities based on a training dataset that includes a plurality of different custom policies deployed under a plurality of different conditions by a plurality of different client entities.

19. The computer program product of claim 16, further comprising computer instructions for:

managing a policy repository configured to store a plurality of community policies that are available for use by a plurality of different client entities; and

storing, in response to the generating of the community policy, the community policy within the policy repository.

* * * * *