

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411634
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Ki; Yang Seok

Parallel processing in computational storage

Abstract

A storage device includes a first storage source storing first data. The storage device further includes a second storage source storing second data. The second data includes a first portion and a second portion separated by a delimiter. The storage device further includes a first buffer configured to receive the first data. The storage device further includes a second buffer configured to receive the second data. The storage device further includes a first processor associated with the first buffer. The storage device further includes a second processor associated with the second buffer. The second processor is configured to perform a first operation on the second portion of the second data, and wherein the first processor is configured to perform a second operation on the first data and the first portion of the second data based on the delimiter.

Inventors:	Ki; Yang Seok (Palo Alto, CA)
Applicant:	Samsung Electronics Co., Ltd. (Suwon-si, KR)
Family ID:	84901767
Assignee:	Samsung Electronics Co., Ltd. (Yongin-si, KR)
Appl. No.:	17/856823
Filed:	July 01, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230280936 A1	Sep. 07, 2023

Related U.S. Application Data

us-provisional-application US 63316307 20220303

Publication Classification

Int. Cl.: G06F3/06 (20060101)

U.S. Cl.:

CPC **G06F3/0656** (20130101); **G06F3/0604** (20130101); **G06F3/065** (20130101);
G06F3/0679 (20130101);

Field of Classification Search

CPC: G06F (3/0656); G06F (3/0604); G06F (3/065); G06F (3/0679); G06F (15/167); G06F
(3/0658); G06F (2212/7203)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6310898	12/2000	Schwartz	N/A	N/A
7441224	12/2007	Bellas et al.	N/A	N/A
7631245	12/2008	Lasser	N/A	N/A
7856544	12/2009	Schenfeld et al.	N/A	N/A
7873779	12/2010	Maddali et al.	N/A	N/A
8149622	12/2011	Lee et al.	N/A	N/A
8261008	12/2011	Que	N/A	N/A
8458377	12/2012	Piccirillo et al.	N/A	N/A
8478736	12/2012	Biran et al.	N/A	N/A
8634247	12/2013	Sprouse et al.	N/A	N/A
8724381	12/2013	Burger, Jr. et al.	N/A	N/A
8984369	12/2014	Varanasi	N/A	N/A
10459661	12/2018	Shivandand et al.	N/A	N/A
10984044	12/2020	Batsakis	N/A	G06F 16/907
11061591	12/2020	Bae et al.	N/A	N/A
11620336	12/2022	Batsakis	707/722	G06F 16/27
2008/0046630	12/2007	Lasser	N/A	N/A
2014/0236889	12/2013	Vasan	707/610	G06F 11/20
2016/0036903	12/2015	Pal	709/213	H04L 67/1087
2016/0191665	12/2015	Kang	709/205	G06F 11/1076
2017/0011022	12/2016	Grossman	N/A	G06F 16/252
2018/0089278	12/2017	Bhattacharjee	N/A	H04L 43/12
2018/0121351	12/2017	Zhang et al.	N/A	N/A
2018/0341549	12/2017	Bolkhovitin	N/A	G06F 11/1068
2019/0370076	12/2018	Behar	N/A	G06F 9/5027
2020/0174701	12/2019	Jung	N/A	G06F 3/061
2020/0249839	12/2019	Yang et al.	N/A	N/A
2020/0257629	12/2019	Pinto	N/A	N/A
2020/0393991	12/2019	Kachare et al.	N/A	N/A
2021/0144428	12/2020	Karazoun	N/A	G06F 16/48
2022/0224624	12/2021	Kasichainula	N/A	H04L 41/0896

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2020/186081	12/2019	WO	N/A

OTHER PUBLICATIONS

EPO Extended European Search Report dated Jul. 18, 2023, issued in corresponding European Patent Application No. 23150952.2 (10 pages). cited by applicant

Wei, et al., "A Hardware-Software Co-design Experiments Platform for NAND Flash Based on Zynq", IEEE Xplore, Library of Congress. Downloaded on Dec. 6, 2021, 7 pgs. cited by applicant

Yoon, et al., Two-Stage In-Storage Processing and Scheduling for Pattern Matching Applications, IEEE Access, <https://creativecommons.org/licenses/by/4.0/>, Creative Commons Attribution 4.0 License, Received May 27, 2021, accepted Jun. 17, 2021, date of publication Jun. 30, 2021, date of current version Jul. 13, 2021, 14 pgs. cited by applicant

Jeong, et al., "REACT: Scalable and High-Performance Regular Expression Pattern Matching Accelerator for In-Storage Processing", IEEE Transactions on Parallel and Distributed Systems, vol. 31, No. 5, May 2020. pp. 1137-1151. cited by applicant

Primary Examiner: Grullon; Francisco A

Attorney, Agent or Firm: Lewis Roca Rothgerber Christie LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATION(S) (1) The present application claims priority to and the benefit of U.S. Provisional Application No. 63/316,307, filed Mar. 3, 2022 entitled "PARALLEL PROCESSING OF STREAM DATA IN COMPUTATIONAL STORAGE," the entire contents of all which is incorporated herein by reference.

FIELD

(1) The present disclosure generally relates to systems and methods for parallel processing in computational storage devices.

BACKGROUND

(2) Computational storage devices (CSD) provide computation functions and data storage. Accordingly, a host may store data at a computational storage device. Host data units (e.g., files) may have a different size than CSD data units (e.g., blocks). Accordingly, a host data unit may be split across multiple components of the CSD.

(3) Computations offloaded to the CSD may be targeted to host data units. While, performing the offloaded computations in parallel may decrease total time spent on the computations, it may be difficult for the CSD to provide a parallel execution pipeline due to unknown alignment of the host data units.

(4) The above information disclosed in this background section is only for enhancement of understanding of the background of the disclosure and therefore it may contain information that does not constitute prior art.

SUMMARY

(5) In various embodiments, described herein include systems, methods, and apparatuses related to resource isolation in computational storage devices.

(6) A storage device includes a first storage source storing first data. The storage device further includes a second storage source storing second data. The second data includes a first portion and a second portion separated by a delimiter. The storage device further includes a first buffer

configured to receive the first data. The storage device further includes a second buffer configured to receive the second data. The storage device further includes a first processor associated with the first buffer. The storage device further includes a second processor associated with the second buffer. The second processor is configured to perform a first operation on the second portion of the second data, and wherein the first processor is configured to perform a second operation on the first data and the first portion of the second data based on the delimiter.

(7) A method includes receiving, at a first buffer, first data from a first storage source. The method further includes receiving, at a second buffer, second data from a second storage source, the second data including a first portion and a second portion separated by a delimiter. The method further includes performing, at a second processor associated with the second buffer, a first operation on the second portion of the second data. The method further includes performing, at a first processor associated with the first buffer, a second operation on the first data and the first portion of the second data.

(8) A storage device includes a first storage channel including a first media device storing first data. The storage device further includes second storage channel including a second media device storing second data. The second data including a first portion and a second portion separated by a delimiter. The storage device further includes a first compute module associated with the first storage channel and including a first processor and a first input buffer. The storage device further includes a second compute module associated with the second storage channel and including a second processor and a second input buffer. The second processor is configured to perform a first operation on the second portion of the second data. The first processor is configured to perform a second operation on the first data and the first portion of the second data based on the delimiter.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The above-mentioned aspects and other aspects of the present techniques will be better understood when the present application is read in view of the following figures in which like numbers indicate similar or identical elements:

(2) FIG. 1 is a block diagram of a system for parallel processing in computational storage.

(3) FIG. 2 is a block diagram illustrating an example of the system in which a processor is configured to transfer data to another processor for processing.

(4) FIG. 3 is a block diagram illustrating an example of the system in which processors associated with storage sources send data from the storage sources to a processor associated with a group of storage sources for processing.

(5) FIG. 4 is a block diagram illustrating an example of the system in which processors read past associated input buffers in a continuous memory space based on delimiter locations.

(6) FIG. 5 is a block diagram of a storage device that supports parallel processing of data based on delimiters.

(7) FIG. 6 is a diagram of a compute module.

(8) FIG. 7 is a diagram of a device performing delimiter based parallel computations in which compute modules operating in parallel pass data between each other.

(9) FIG. 8 showing examples of input buffer structures.

(10) FIG. 9 is a diagram showing a process 900 for performing delimiter based parallel processing.

(11) FIG. 10 is a diagram illustrating an example of a device performing delimiter based parallel computations in which compute modules operating in parallel pass data to one or more downstream processors based on delimiters.

(12) FIG. 11 is a diagram illustrating an example of a device performing delimiter based parallel computations in which compute modules operating in parallel directly access input buffers of other

compute modules based on delimiter locations.

(13) FIG. 12 is a flowchart of a method for delimiter based parallel processing.

(14) While the present techniques are susceptible to various modifications and alternative forms, specific embodiments thereof are shown by way of example in the drawings and will herein be described. The drawings may not be to scale. It should be understood, however, that the drawings and detailed description thereto are not intended to limit the present techniques to the particular form disclosed, but to the contrary, the intention is to cover all modifications, equivalents, and alternatives falling within the spirit and scope of the present techniques as defined by the appended claims.

DETAILED DESCRIPTION OF DRAWINGS

(15) The details of one or more embodiments of the subject matter described herein are set forth in the accompanying drawings and the description below. Other features, aspects, and advantages of the subject matter will become apparent from the description, the drawings, and the claims.

(16) Various embodiments of the present disclosure now will be described more fully hereinafter with reference to the accompanying drawings, in which some, but not all embodiments are shown. Indeed, the disclosure may be embodied in many different forms and should not be construed as limited to the embodiments set forth herein; rather, these embodiments are provided so that this disclosure will satisfy applicable legal requirements. The term “or” is used herein in both the alternative and conjunctive sense, unless otherwise indicated. The terms “illustrative” and “example” are used to be examples with no indication of quality level. Like numbers refer to like elements throughout. Arrows in each of the figures depict bi-directional data flow and/or bi-directional data flow capabilities. The terms “path,” “pathway” and “route” are used interchangeably herein.

(17) Embodiments of the present disclosure may be implemented in various ways, including as computer program products that comprise articles of manufacture. A computer program product may include a non-transitory computer-readable storage medium storing applications, programs, program components, scripts, source code, program code, object code, byte code, compiled code, interpreted code, machine code, executable instructions, and/or the like (also referred to herein as executable instructions, instructions for execution, computer program products, program code, and/or similar terms used herein interchangeably). Such non-transitory computer-readable storage media include all computer-readable media (including volatile and non-volatile media).

(18) In one embodiment, a non-volatile computer-readable storage medium may include a floppy disk, flexible disk, hard disk, solid-state storage (SSS) (for example a solid-state drive (SSD)), solid state card (SSC), solid state component (SSM), enterprise flash drive, magnetic tape, or any other non-transitory magnetic medium, and/or the like. A non-volatile computer-readable storage medium may also include a punch card, paper tape, optical mark sheet (or any other physical medium with patterns of holes or other optically recognizable indicia), compact disc read only memory (CD-ROM), compact disc-rewritable (CD-RW), digital versatile disc (DVD), Blu-ray disc (BD), any other non-transitory optical medium, and/or the like. Such a non-volatile computer-readable storage medium may also include read-only memory (ROM), programmable read-only memory (PROM), erasable programmable read-only memory (EPROM), electrically erasable programmable read-only memory (EEPROM), flash memory (for example Serial, NAND, NOR, and/or the like), multimedia memory cards (MMC), secure digital (SD) memory cards, SmartMedia cards, CompactFlash (CF) cards, Memory Sticks, and/or the like. Further, a non-volatile computer-readable storage medium may also include conductive-bridging random access memory (CBRAM), phase-change random access memory (PRAM), ferroelectric random-access memory (FeRAM), non-volatile random-access memory (NVRAM), magnetoresistive random-access memory (MRAM), resistive random-access memory (RRAM), Silicon-Oxide-Nitride-Oxide-Silicon memory (SONOS), floating junction gate random access memory (FJG RAM), Millipede memory, racetrack memory, and/or the like.

(19) In one embodiment, a volatile computer-readable storage medium may include random access memory (RAM), dynamic random access memory (DRAM), static random access memory (SRAM), fast page mode dynamic random access memory (FPM DRAM), extended data-out dynamic random access memory (EDO DRAM), synchronous dynamic random access memory (SDRAM), double data rate synchronous dynamic random access memory (DDR SDRAM), double data rate type two synchronous dynamic random access memory (DDR2 SDRAM), double data rate type three synchronous dynamic random access memory (DDR3 SDRAM), Rambus dynamic random access memory (RDRAM), Twin Transistor RAM (TTRAM), Thyristor RAM (T-RAM), Zero-capacitor (Z-RAM), Rambus in-line memory component (RIMM), dual in-line memory component (DIMM), single in-line memory component (SIMM), video random access memory (VRAM), cache memory (including various levels), flash memory, register memory, and/or the like. It will be appreciated that where embodiments are described to use a computer-readable storage medium, other types of computer-readable storage media may be substituted for or used in addition to the computer-readable storage media described above.

(20) As should be appreciated, various embodiments of the present disclosure may also be implemented as methods, apparatus, systems, computing devices, computing entities, and/or the like. As such, embodiments of the present disclosure may take the form of an apparatus, system, computing device, computing entity, and/or the like executing instructions stored on a computer-readable storage medium to perform certain steps or operations. Thus, embodiments of the present disclosure may also take the form of an entirely hardware embodiment, an entirely computer program product embodiment, and/or an embodiment that comprises combination of computer program products and hardware performing certain steps or operations.

(21) Embodiments of the present disclosure are described below with reference to block diagrams and flowchart illustrations. Thus, it should be understood that each block of the block diagrams and flowchart illustrations may be implemented in the form of a computer program product, an entirely hardware embodiment, a combination of hardware and computer program products, and/or apparatus, systems, computing devices, computing entities, and/or the like carrying out instructions, operations, steps, and similar words used interchangeably (for example the executable instructions, instructions for execution, program code, and/or the like) on a computer-readable storage medium for execution. For example, retrieval, loading, and execution of code may be performed sequentially such that one instruction is retrieved, loaded, and executed at a time. In some example embodiments, retrieval, loading, and/or execution may be performed in parallel such that multiple instructions are retrieved, loaded, and/or executed together. Thus, such embodiments can produce specifically-configured machines performing the steps or operations specified in the block diagrams and flowchart illustrations. Accordingly, the block diagrams and flowchart illustrations support various combinations of embodiments for performing the specified instructions, operations, or steps.

(22) As used herein, a computational storage device (CSD) refers to a storage device that supports computational tasks. For example, a CSD may include a storage element (e.g., non-volatile memory, such as flash memory, a hard disk drive, etc.) and a compute element (e.g., a central processor unit (CPU), graphics processor unit (GPU), a field programmable gate array (FPGA), an application specific integrated circuit (ASIC) (such as a tensor processing unit), processor core, etc.) and be configured to support storage of data at the compute element and execution of computational tasks at the compute element. Accordingly, a CSD may provide storage capabilities to a host device (e.g., a computing device) and may support offloading of computational tasks from the host device to the CSD device.

(23) In some examples according to the disclosure, a computational storage device (CSD) includes more than one compute engine and more than one storage source. Examples of media sources include, storage media (e.g., flash chips, such as NAND flash chips, flash media packages, resistive random access memory devices, hard disk devices, etc.) storage channels (e.g., NAND flash

channels, etc.), other groupings of storage media, etc. The compute engines receive data from the storage sources and perform computations on the data. Because host data units (e.g., files) operated on by the computations may be split across more than one storage source, the compute engines operate on the data based on locations of delimiters that indicate boundaries between host data units in the data. In particular, a compute engine may begin performing computations on data that follows a first instance of a delimiter in a buffer of the compute engine. Data that precedes the first delimiter may be combined with data from an input buffer of a previous compute engine and processed elsewhere. Similarly, a compute engine may detect a final instance of a delimiter in the input buffer associated with the compute engine and stop computations at the final instance until additional data is available. In some examples, data following the final instance of the delimiter may be carried over to another input buffer to be processed by another compute engine.

(24) The disclosed delimiter aware systems and methods may provide for parallel computations in a CSD despite misalignment between host data units and CSD data units. These systems and methods may be particularly useful in RAID configurations in which data is striped across several storage sources. Further, the disclosure may be extended to systems that include host data stored across more than one CSD.

(25) Referring to FIG. 1, a block diagram of a system **100** for parallel processing in computational storage is shown. The system **100** includes a first processor **102**, a second processor **104**, a first buffer **106**, a second buffer **108**, a first storage source **110**, and a second storage source **112**. The first processor **102** may include a central processor unit, a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), a graphics processor unit (GPU), another type of processor, or any combination thereof. The second processor **104** may similarly include any of the types of processor described with respect to the first processor **102**. In some implementations the first processor **102**, the second processor **104**, or both may include a portion of a processor device. For example, the first processor **102** and the second processor **104** may each correspond to 1 or more processor cores.

(26) The first buffer **106** may include a memory device, such as a dynamic random access memory (DRAM), a static random access memory (SRAM), another type of memory, or a combination thereof. The second buffer **108** may similarly include any of the types of memory device described with respect to the first buffer **106**. In some implementations, the first buffer **106** and the second buffer **108** correspond to different regions of the same memory device (or virtual memory device).

(27) The first storage source **110** includes one or more storage devices, such as a flash chip (e.g., NAND flash), a flash package, a flash channel (e.g., a NAND flash channel), a hard disk device (HDD), a resistive random access memory (RRAM) device, etc. The second storage source **112** may similarly include any of the types of storage device described with respect to the first storage source **110**.

(28) The first processor **102** is associated with the first buffer **106**. In some implementations, the first buffer **106** is included within the first processor **102** or the first processor **102** and the first buffer **106** are included in a common compute module. In other implementations, the first processor **102** and first buffer **106** are distinct components of the system **100**, and the first processor **102** is configured to utilize the first buffer **106** as an input buffer.

(29) The second processor **104** is associated with the second buffer **108**. In some implementations, the second buffer **108** is included within the second processor **104** or the second processor **104** and the second buffer **108** are included in a common compute module. In other implementations, the second processor **104** and second buffer **108** are distinct components of the system **100**, and the second processor **104** is configured to utilize the second buffer **108** as an input buffer.

(30) The first storage source **110** and the second storage source **112** may store host data. Because host data units (e.g., files) may not be aligned with storage source data units (e.g., chunks), a host data unit may be split across the first storage source **110** and the second storage source **112**.

Individual host data units may be separated by delimiters. Accordingly, data occurring between 2

consecutive delimiters may correspond to a single host data unit. Computations to be performed by the first processor **102** and the second processor **104** may be targeted to whole host data units.

(31) In the illustrated example, the first storage source **110** stores first data (e.g., first host data) including a first delimiter **113** and a first portion **114**. The second storage source **112** stores second data (e.g., second host data) including a second portion **116**, a second delimiter **118**, a third portion **120**, and a third delimiter **121**. The first portion **114** and the second portion **116** may comprise a single host data unit. Accordingly, successful computation based on the single data unit may be based on a processor having access to both the first portion **114** and the second portion **116**.

(32) In operation, the first delimiter **113** and the first data **114** are transferred (e.g., copied) from the first storage source **110** into the first buffer **106**, and the second portion **116**, the second delimiter **118**, the third portion **120**, and the third delimiter **121** are transferred (e.g., copied) from the second storage source **112** to the second buffer **108**. For example, the transfers may correspond to direct memory access (DMA) transfers or some other type of transfers. Transfer of the first delimiter **113** and the first portion **114** may be initiated by the first processor **102**. Transfer of the second portion **116**, the second delimiter **118**, the third portion **120**, and the third delimiter **121** of the second data may be initiated by the second processor **104**. The first delimiter **113**, the first portion **114**, the second portion **116**, the second delimiter **118**, the third portion **120**, and the third delimiter **121** may be transferred in response to a request (e.g., from a host device) to perform one or more computations based on host data stored in the first storage source **110** and the second storage source **112**. As described above, a computation may target a whole host data unit.

(33) The first processor **102** is configured to identify locations of delimiters in data transferred into the first buffer **106**. These delimiters may separate host data units from each other. Similarly, the second processor **104** is configured to identify locations of delimiters in data transferred into the second buffer **108**. Based on locations of the delimiters, the first processor **102** determines which data in the first buffer **106** to process. Similarly, the second processor **104** determines which data in the second buffer **108** to process based on locations of delimiters in the second buffer **108**. Further, the first processor **102** and/or the second processor **104** may perform computations on data from a different input buffer (e.g., a buffer associated with a different processor) based on locations of delimiters in data of the different input buffer.

(34) For example, the second processor **104** (or the first processor **102**) may be configured to ignore data that appears before a first delimiter in the second buffer **108** (or the first buffer **106**). Such data may be processed by a different processor. To illustrate, data that appears before the first delimiter may be correspond to an incomplete host data unit. Accordingly, a processor (e.g., the first processor **102**) that has the rest of the host data unit may process this data. The first processor **102** may handle data that occurs before a first delimiter in the first buffer similarly.

(35) Further, the second processor **104** may ignore data that appears after a last delimiter in the second buffer **108** (or the first buffer **106**). This data that follows the last delimiter may be processed by the second processor **104** in combination with additional data from another input buffer or may be processed in combination with the additional data by a different processor. For example, the second processor **104** may wait until additional data is passed into the second buffer **108** to complete a host data unit following the last delimiter. The first processor **102** may handle data that follows a last delimiter in the first buffer **106**.

(36) Data that occurs in between the first and last delimiter in the second buffer **108** may be processed by the second processor **104**. Similarly, data that occurs between a first and last delimiter in the first buffer **106** may be processed by the first processor **102**. By selectively performing computations based on data in the second buffer **108** according to first and last delimiter locations, the second processor **104** may perform computations on whole host data units.

(37) In the illustrated example, the second processor **104** identifies a location of the second delimiter **118** and a location of the third delimiter **121** in the second buffer **108**. Based on the location of the second delimiter **118** (e.g., first delimiter in the second buffer **108**), the second

processor **104** refrains from processing the second portion **116**. “Refraining” from processing the second portion **116** may include beginning to perform computations on data stored at a location in the buffer after the second delimiter **118**, transferring (e.g., copying) the second portion **116** to an input buffer of another processor, or a combination thereof.

(38) The second processor **104** performs computations on the third portion **120**, as indicated by arrow **126**, based on the location of the second delimiter **118** (e.g., the second delimiter **118**) and the location of the third delimiter **121**. For example, because the third portion **120** falling between a first and last delimiter in the second buffer **108**, the second processor **104** may perform a computation on the third portion **120**. On the other hand, the first processor **102** performs a computation based on a combination of the first portion **114** and the second portion **116**, as indicated by arrows **122** and **124**, based on a position of the first delimiter **113** and the second delimiter **118**. In some implementations, the second processor **104** transfers the second portions **116** into the first buffer **106**. In other implementations, the first buffer **106** and the second buffer **108** are arranged in a continuous memory space and the first processor **102** continues performing computations based a boundary of the first buffer **106** until the second delimiter **118** is reached.

(39) Thus, in the illustrated example, a host data unit comprising the first portion **114** and the second portion **116** may be processed by the same processor (e.g., the first processor **102**) despite originating from distance storage sources and being initially copied into input buffers of different processors. Accordingly, the system **100** may provide parallel processing of host data despite host data units being unaligned with data units of the storage system **100** and falling across several storage sources of the system **100**. Therefore, the system **100** may be suited for various implementations in which host data may be broken across several storage sources, such as RAID systems.

(40) The system **100** may have alternative configurations than illustrated in FIG. 1. For example, the system **100** may include a different number of processors (and corresponding buffers). Further, while the first buffer **106** and the second buffer **108** are illustrated as distinct components, they may be portions of a continuous memory space. Further, the system **100** may include a different number of storage sources. Additionally, components other than those shown may be included in the system **100**. It should also be noted that components of the system **100** may be virtual components provided by the system **100** (e.g., by a processor executing a hypervisor or other emulation software). Further, while FIG. 1 illustrates a 1-to-1 correspondence between processors and storage sources, it should be noted that in some implementations a processor may be configured to receive and process data from more than one storage source. In some of these examples, a processor may be associated with (e.g., process data from) each memory source and then an additional processor may be associated with groups of memory sources. For example, each flash chip may have an associated processor and each flash channel (including several flash chips) may have an associated processor.

(41) FIG. 2 is a block diagram illustrating an example of the system **100** in which a processor is configured to transfer data to another processor for processing. In the illustrated example, the second processor **104** transfers (e.g., copies) the second portion **116** to the first buffer **106** in response to the second portion **116** being located in front of a first delimiter (e.g., the second delimiter **118**) in the second buffer **108**, as indicated by arrow **222**. The second processor **104** may utilize a DMA operation to transfer the second portion **116**. In some implementations, the second processor **104** further transfers the second delimiter **118** to the first buffer **106**. Transferring the second portion **116** to the first buffer **106** unites a host data unit comprising the first portion **114** and the second portion **116**. Accordingly, the first processor **102** may perform a computation on the host data unit comprising the first portion **114** and the second portion **116**.

(42) FIG. 3 is a block diagram illustrating an example of the system **100** in which processors associated with storage sources send data from the storage sources to a processor associated with a group of storage sources for processing. In the illustrated example, a third processor **302** is directly

associated with the first storage source **110** and the first processor **102** is associated with a group including the first storage source **110** and the second storage source **112**. For example, the third processor **302** may be configured to process data from a first flash chip corresponding to the first storage source **110**, the second storage source **112** may be configured to process data from a second flash chip corresponding to the second storage source **112**, and the first processor **102** may be configured to process data from a flash channel including the first storage source **110** and the second storage source **112**.

(43) In the illustrated example, a fourth delimiter **308**, a fourth portion **306** of data, the first delimiter **113**, and the first portion **114** are transferred (e.g., via DMA initiated by the third processor **302**) from the first storage source **110** to a third buffer **304** of the third processor **302**. The third processor **302** processes the fourth portion **306** based on the fourth portion **306** falling between a first delimiter (e.g., the fourth delimiter **308**) in the third buffer **304** and a last delimiter (e.g., the first delimiter **113**) in the third buffer **304**. Based on the first portion **114** being located after the last delimiter in the third buffer **304**, the third processor **302** initiates transfer of the first portion **114** to the first buffer **106**, as indicated by arrow **322**. In some examples, the third processor **302** transfers the first delimiter **113** in addition to the first portion **114**. Further, based on the second portion **116** being located before a first delimiter (e.g., the second delimiter **118**) in the second buffer, the second processor **104** initiates transfer of the second portion **116** to the first buffer **106**, as indicated by arrow **324**. In some examples, the second processor **104** transfers the second delimiter **118** in addition to the second portion **116**. The first processor **102** may perform a computation on a whole host data unit comprising the first portion **114** and the second portion **116**.

(44) FIG. 4 is a block diagram illustrating an example of the system **100** in which processors read past associated input buffers in a continuous memory space based on delimiter locations. In the example of FIG. 4, the first buffer **106** and the second buffer **108** are part of a common buffer space **402**. The buffer space **402** may correspond to a single physical memory device or to a virtual memory space supported by several memory devices. The first buffer **106** corresponds to an input buffer for the first processor **102** and the second buffer **108** corresponds to an input buffer for the second processor **104**. However, the first processor **102** (and the second processor **104**) may ignore data in an associated input buffer and/or process data in an input buffer of another processor based on delimiter locations. Ignoring data may include beginning processing at a location after that data.

(45) In the illustrated example, the first processor **102** begins performing computations on data in the first buffer after a first delimiter (e.g., the first delimiter **113**) in the first buffer **106**, as indicated by arrow **422**. In response to detecting data following a last delimiter (e.g., the first delimiter **113**) in the first buffer **106**, the first processor **102** further continues processing until reaching a first delimiter (e.g., the second delimiter **118**) in a next input buffer (e.g., the second buffer **108**), as indicated by arrow **423**. Accordingly, the first processor **102** may perform one or more computations based on an entire host data unit comprising the first portion **114** and the second portion **116** despite the first portion **114** and the second portion originating from different storage sources and being input to different input buffers associated with different processors. Further, the second processor **104** may begin performing computations on data in the second buffer **108** that occurs following a first delimiter (e.g., the second delimiter **118**) in the second buffer **108**, as indicated by arrow **426**. The second processor **104** continues until reaching a final delimiter (e.g., the third delimiter **121**) in the second buffer **108**, as shown by arrow **427**. In examples in which additional data follows the final delimiter in the second buffer **108**, the second processor **104** may continue to perform computations on data in a next buffer (not shown).

(46) As illustrated, data between host data unit delimiters is processed by a single processor. Thus, FIG. 4 represents another example of how the system may perform parallel processing on data despite misalignment between host data units and storage system data units.

(47) Referring to FIG. 5, an example of a storage device **500** that supports parallel processing of data is shown. The storage device **500** is a computational storage device and may include a solid

state storage device (SSD), a hard disk device, another type of storage device, or a combination thereof. In the illustrated example, the storage device **500** includes a storage device controller **504** and storage media. The storage media includes a first channel **534** (e.g., a flash channel or other type of storage channel) that includes media device (e.g., a storage media device) **534a**, media device **534b**, media device **534c**, and media device **534d**. The storage media further includes a second channel **536** (e.g., a flash channel or other type of storage channel) that includes media device **536a**, media device **536b**, media device **536c**, and media device **536d**. The storage media further includes a third channel **538** (e.g., a flash channel or other type of storage channel) that includes media device **538a**, media device **538b**, media device **538c**, and media device **538d**. The storage media further includes a fourth channel **540** (e.g., a flash channel or other type of storage channel) that includes media device **540a**, media device **540b**, media device **540c**, and media device **540d**. The media devices **534a-d**, **536a-d**, **538a-d**, **540a-d** may include NAND flash chips, hard disk media, other types of storage media, or a combination thereof.

(48) The first storage source **110** and the second storage source **112** of FIGS. **1-4** may correspond to different elements of a group including the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, one of the channels **534**, **536**, **538**, **540**, a group of the channels **534**, **536**, **538**, **540**, or some other aspect of the storage media of the storage device **500**. In some examples, the storage media has a different configuration than shown. For example, the storage device **500** may include more or fewer channels, more or fewer media devices per channel, more or fewer channels per channel group, or a combination thereof. Further, it should be noted that in some implementations, elements of the storage device **500** are virtual.

(49) The storage device **500** further includes a storage device memory **512**. The storage device memory **512** may include a memory device, such as a DRAM device, an SRAM device, another type of memory device, or a combination thereof. In some examples, the storage device memory **512** includes more than one device.

(50) The storage device controller **504** includes a host interface controller **506**, a controller memory **505**, a storage device compute module **508**, a memory controller **510**, a first storage media controller **514**, and a second storage media controller **522**. The host interface controller **506** may include hardware components, firmware, software, or a combination thereof configured to provide an interface to a host device. In some implementations, the interface includes a non-volatile memory express (NVMe) interface, a compute express link (CXL) interface, or another type of storage interface.

(51) The controller memory **505** may include a memory device, such as a DRAM device, an SRAM device, another type of memory device, or a combination thereof. In some examples, the controller memory **505** includes more than one device. In some examples, the controller memory **505** and the storage device memory **512** are different types of memory with different characteristics (e.g., latency, capacity, etc.). To illustrate, the controller memory **505** may include SRAM while the storage device memory **512** includes DRAM.

(52) The storage device controller **504** further includes a storage device compute module **508**. The storage device controller **504** is configured to perform computations based on data stored in the storage media of the storage device **500**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the storage device compute module **508** may include an input buffer, an output buffer, and a processor. The input and output buffers may be separate components included in the storage device compute module **508** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(53) The first storage media controller **514** includes a flash controller or other type of storage controller. In various examples, the first storage media controller **514** includes hardware, firmware, software, or a combination thereof configured to control access to the first channel **534** (and associated media devices **534a-d**, **536a-d**) and the second channel **536** (and associated media

devices **536a-d**, **536a-d**). Controlling access may include performing translations of memory addresses to and from an address space used by the media devices **534a-d**, **536a-d** and initiating data transfers to and from the media devices **534a-d**, **536a-d**.

(54) Similarly, the second storage media controller **522** includes a flash controller or other type of storage controller. In various examples, the second storage media controller **522** includes hardware, firmware, software, or a combination thereof configured to control access to the third channel **538** (and associated media devices **538a-d**, **538a-d**) and the fourth channel **540** (and associated media devices **540a-d**, **540a-d**). Controlling access may include performing translations of memory addresses to and from an address space used by the media devices **538a-d**, **538a-d** and initiating data transfers to and from the media devices **540a-d**, **540a-d**.

(55) The first storage media controller **514** further includes a first media core compute module **516**. The first media core compute module **516** is configured to perform computations based on data stored in the media devices **534a-d**, **536a-d** connected to the first storage media controller **514**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the first media core compute module **516** may include an input buffer, an output buffer, and a processor. The input and output buffers may be distinct components included in the first media core compute module **516** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(56) The first storage media controller **514** further includes a first channel compute module **518**. The first channel compute module **518** is configured to perform computations based on data stored in the media devices **534a-d** of the first channel **534**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the first channel compute module **518** may include an input buffer, an output buffer, and a processor. The input and output buffers may be distinct components included in the first channel compute module **518** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(57) The first storage media controller **514** further includes a second channel compute module **520** configured to perform computations based on data stored in the media devices **536a-d** of the second channel **536**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the second channel compute module **520** may include an input buffer, an output buffer, and a processor. The input and output buffers may be distinct components included in the second channel compute module **520** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(58) The second storage media controller **522** further includes a second media core compute module **524**. The second media core compute module **524** is configured to perform computations based on data stored in the media devices **538a-d**, **540a-d** connected to the second storage media controller **522**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the second media core compute module **524** may include an input buffer, an output buffer, and a processor. The input and output buffers may be distinct components included in the second media core compute module **524** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(59) The second storage media controller **522** further includes a third channel compute module **526**. The third channel compute module **526** is configured to perform computations based on data stored in the media devices **538a-d** of the third channel **538**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the third channel compute module **526** may include an input buffer, an output buffer, and a processor. The input and output buffers may be distinct components included in the third channel compute module **526** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(60) The first storage media controller **514** further includes a second channel compute module **520**

configured to perform computations based on data stored in the media devices **536a-d** of the second channel **536**. Such computations may include filter operations, mathematical operations, search operations, etc. As described further herein, the second channel compute module **520** may include an input buffer, an output buffer, and a processor. The input and output buffers may be separate components included in the first storage media controller **514** or may include ranges of memory included in the controller memory **505** and/or the storage device memory **512**.

(61) It should be noted that the compute modules **516**, **518**, **520**, **524**, **526**, **528** may be external to the storage media controllers **514**, **522** in some implementations. Further, the storage device **500** may include a different number of channel compute modules and/or media core compute modules than shown. While not illustrated, one or more of the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d** may have an associated compute module. These compute modules may have a structure similar to those of the channel compute modules **518**, **520**, **526**, **528** and the media core compute modules **516**, **524** and each may be configured to perform computations on data from a respective media device. In some implementations, the first processor **102** and the first buffer **106** correspond to one of the compute modules of the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, the channel compute modules **518**, **520**, **526**, **528**, and the media core compute modules **516**, **524**. The second processor **104** and the second buffer **108** may correspond to a different one of the compute modules of the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, the channel compute modules **518**, **520**, **526**, **528**, and the media core compute modules **516**, **524**. The third processor **302** and the third buffer **304** may correspond to one of the channel compute modules **518**, **520**, **526**, **528**, the media core compute modules **516**, **524**, and the storage device compute module **508**.

(62) Each of the first buffer **106**, the second buffer **108**, and the third buffer **304** may be implemented in the controller memory **505**, in the storage device memory **512**, in a memory of a compute module corresponding to a media device, in a memory of the first channel compute module **518**, in a memory of the second channel compute module **520**, in a memory of the third channel compute module **526**, in a memory of the fourth channel compute module **528**, in a memory of the first media core compute module **516**, in a memory of the second media core compute module **524**, in a memory of the storage device compute module **508**, in a different component of the storage device **500**, or a combination thereof.

(63) In response to a host command (or other trigger), the storage device controller **504** may initiate transfer of data from one or more of the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d** to one or more input buffers of compute modules. The compute modules may include compute modules associated with the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, the channel compute modules **518**, **520**, **526**, **528**, the media core compute modules **516**, **524**, the storage device compute module **508**, or a combination thereof. The compute modules that receive the data perform one or more computations on the data to generate results. However, individual host data units may be split across input buffers because the individual host data units may be split across the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, the channel compute modules **518**, **520**, **526**, **528**. In order to process whole host data units, the compute modules perform the computations based on locations of delimiters in the input buffers, as described herein. For example, based on locations of delimiters in the data, compute modules may transfer data to other input buffers and/or may access input buffers associated with other compute modules. Additionally, in some implementations, compute modules of the storage device **500** are configured to operate on data as the data is streamed into (e.g., read into) the corresponding input buffer. Based on delimiter position in an input buffer, a compute module may delay performing computations until an entire host data unit is available to the compute module. Accordingly, an individual host data unit may be processed by a single processor.

(64) Referring to FIG. 6, an example of a compute module **600** is shown. Compute modules of the media devices **534a-d**, **536a-d**, **538a-d**, **540a-d**, the first media core compute module **516**, the first channel compute module **518**, the second channel compute module **520**, the second media core

compute module **524**, the third channel compute module **526**, the fourth channel compute module **528**, and the storage device compute module **508** may be analogous to the compute module **600**. The compute module **600** includes an output buffer **604**, a processor **606**, and an input buffer **608**. The output buffer **604** may correspond to a memory device such, as a DRAM or SRAM device. Similarly, the input buffer **608** may correspond to a memory device, such as a DRAM or SRAM device. In some implementations, the output buffer **604** and the input buffer **608** correspond to regions of one memory (e.g., a physical or virtual memory space). While illustrated as components of the compute module **600**, it should be noted that the output buffer **604** and the input buffer **608** may be implemented as regions of a memory external to the compute module **600**, such as the storage device memory **512**, the controller memory **505**, or a combination thereof. The input buffer **608** may correspond to the first buffer **106**, the second buffer **108**, or the third buffer **304**.

(65) The processor **606** may correspond to the first processor **102**, the second processor **104**, or the third processor **302**. The processor **606** may include a FPGA, a CPU, a GPU, an ASIC, another type of processor, or a combination thereof. The processor **606** is configured to perform computations based on data in the input buffer **608** and output results of the computations to the output buffer **604**. As described above and below, the processor **606** may perform the computations based on locations of delimiters in the input buffer **608**. Based on these delimiters, the processor **606** may transfer (e.g., copy) data to another input buffer for processing by another processor, process data outside of the input buffer **608**, or a combination thereof in order to keep host data units together. Additionally, the processor **606** may delay performing computations based on delimiters in the input buffer **608**.

(66) Storage devices (e.g., the storage device **500**) and storage systems (e.g., the system **100**) disclosed herein perform delimiter based processing of data stored in storage media. In some implementations, compute modules operating in parallel pass data between each other based on the delimiters. In some implementations, a compute module passes to another compute module downstream in a parallel computing hierarchy based on the delimiters. In some implementations, a compute module reads past a boundary of an associated input buffer into another compute module's input buffer based on the delimiters. Aspects of these implementations may be combined.

(67) FIG. 7 is a diagram illustrating an example of a device **700** performing delimiter based parallel computations in which compute modules operating in parallel pass data between each other. The device **700** may correspond to the storage device **500** of FIG. 5. FIG. 7 depicts a first processor **718**, a second processor **720**, a third processor **722**, and a fourth processor **724**. The first processor **718** has a corresponding first input buffer **710**, the second processor **720** has a corresponding second input buffer **712**, the third processor **722** has a corresponding third input buffer **714**, and the fourth processor **724** has a corresponding fourth input buffer **716**. Further, the first processor **718** has a corresponding first output buffer **758**, the second processor **720** has a corresponding second output buffer **760**, the third processor has a corresponding third output buffer **762**, and the fourth processor has a corresponding fourth output buffer **764**. Each processor and corresponding input buffer and output buffer may be a compute module analogous to the compute module **600**.

(68) The first input buffer **710** is configured to receive data from a first storage source **702**, the second input buffer **712** is configured to receive data from a second storage source **704**, the third input buffer **714** is configured to receive data from a third storage source **706**, and the fourth input buffer **716** is configured to receive data from a fourth storage source **708**. In the illustrated example, the storage sources correspond to storage channels. Accordingly, the first storage source **702** may correspond to the first channel **534**, the second storage source **704** may correspond to the second channel **536**, the third storage source **706** may correspond to the third channel **538**, and the fourth storage source **708** may correspond to the fourth channel **540**. Further, the first input buffer **710**, the first processor **718**, and the first output buffer **758** may correspond to the first channel compute module **518**; the second input buffer **712**, the second processor **720**, and the second output buffer **760** may correspond to the second channel compute module **520**; the third input buffer **714**,

the third processor **722**, and the third output buffer **762** may correspond to the third channel compute module **526**; and the fourth input buffer **716**, the fourth processor **724**, and the fourth output buffer **764** may correspond to the fourth channel compute module **528**.

(69) The device **700** further includes a gathered output buffer **770**. The gathered output buffer **770** may correspond to an output buffer of a compute module downstream from the processors **718**, **720**, **722**, **724**. As used herein, “downstream” indicates a direction toward an output of a processing pipeline. For example, the gathered output buffer **770** may correspond to an output buffer of the storage device compute module **508**, which is downstream of the channel compute modules **518**, **520**, **526**, **528**.

(70) In operation, data is copied from the first storage source **702** to the first input buffer **710**, data is copied from the second storage source **704** to the second input buffer **712**, data is copied from the third storage source **706** to the third input buffer **714**, and data is copied from the fourth storage source **708** to the fourth input buffer **716**. The data may be copied into the input buffers **710**, **712**, **714**, **716** using a DMA transfer or other type of memory transfer. In some implementations, data copy operations are initiated in response to a command to perform an operation on the data (e.g., a search operation, a filter operation, a mathematical operation, etc.). Such a command may be received from a host device (e.g., through the host interface controller **506**). The data copy operations may be initiated by the processors **718**, **720**, **722**, **724**, by an upstream processor, or by another component of a storage device controller (e.g., the storage device controller **504**).

(71) Data is copied into the input buffers **710**, **712**, **714**, **716** in sequence such that a whole host data unit is defined by two adjacent delimiters. Further, data copied into a final position of an input buffer is succeeded in the sequence by data copied into a first position of a next input buffer. Thus, data received by the first input buffer **710** may be followed by data received by the second input buffer **712** in the sequence. Similarly, data received by the second input buffer **712** may be followed by data received by the third input buffer **714** in the sequence and data received by the third input buffer **714** may be followed by data received by the fourth input buffer **716** in the sequence. Data may be streamed into the input buffers **710**, **712**, **714**, **716** as computations are performed by the processors **718**, **720**, **722**, **724**.

(72) The first input buffer **710** (e.g., the first input buffer in a parallel processing pipeline stage) may include additional space at a beginning and an end of the input buffer. Data copied from the first storage source **702** may be placed between the additional space at the beginning and the additional space at the end. The additional space may correspond to host data unit size (e.g., a file size), to a device data unit size (e.g., a block), to another unit, or to a multiple of any of these data unit sizes.

(73) The fourth input buffer **716** (e.g., the last input buffer in a parallel processing pipeline stage) may not include additional buffer space. The second input buffer **712** and the third input buffer **714** (e.g., input buffers between the first and last input buffers in a parallel processing stage) may each include additional space at ends of the input buffers **712**, **714**.

(74) As data is streamed copied into the first input buffer **718**, the first processor **718** is configured to track a position of a first delimiter in the first input buffer **710** and a position of a last delimiter in the first input buffer **710**. Similarly, the second processor **720** tracks positions of first and last delimiters in the second input buffer **712**, the third processor **722** tracks positions of first and last delimiters in the third input buffer **714**, and the fourth processor **724** tracks positions of first and last delimiters in the fourth input buffer **716**. The processors **718**, **720**, **722**, **724** may perform computations on data falling between the first and last delimiters in the respective input buffers.

(75) The second processor **720**, the third processor **722**, and the fourth processor **724** (e.g., processors after the first processor the first processor **718** in a parallel processing stage) are configured to transfer data occurring in a respective input buffer before the first delimiter to a previous input buffer, as part of a forwarding process. In FIG. 7, the second processor **720** initiates transfer of data occurring before the first delimiter in the second input buffer **712** into the additional

space at the end of the first input buffer **710**. Similarly, the third processor **722** initiates transfer of data occurring before the first delimiter in the third input buffer to the additional space at the end of the second input buffer **712**, and the fourth processor **724** initiates transfer of data occurring before the first delimiter in the fourth input buffer **716** to the additional space at the end of the third input buffer **714**. The first processor **718** may perform computations on data transferred into the first input buffer **710** along with data following the last delimiter in the first input buffer **710**. Further, the second processor **720** may perform computations on data transferred into the second input buffer **712** by the third processor **722** along with data following the last delimiter in the second input buffer **712**. Further, the third processor **722** may perform computations on data transferred into the third input buffer **714** by the fourth processor **724** along with data following the last delimiter in the third input buffer **714**. Accordingly, a portion of a host data unit that falls before a first delimiter in an input buffer may be reunited with a remaining portion of the host data unit in a previous input buffer prior to processing. In some implementations, the processors **720**, **722**, **724** transfer the first delimiter to the previous input buffer in addition to transferring data occurring before the first delimiter.

(76) The fourth processor **724** (e.g., a final processor in a parallel processing stage) further transfers data occurring after a final delimiter in the fourth input buffer **716** to the additional space at the beginning of the first input buffer **710** (e.g., the first input buffer in the parallel processing stage) as part of a carryover process. The first processor **718** may perform computations on the data transferred into the first input buffer **710** by the fourth processor **724** with data that occurs before the first delimiter of the first input buffer **710**. Accordingly, a host data unit split between an end of the fourth input buffer **716** and the beginning of the first input buffer **710** may be reunited prior to processing. The final processor may further transfer the final delimiter with the data.

(77) Each of the processors **718**, **720**, **722**, **724** in the device **700** includes four kernels. A different number of kernels may be included than shown. Each kernel may include hardware and/or executable software configured to perform a computation on data (e.g., a search function, a filter function, a mathematical operation, etc.) to generate output. The first processor **718** includes a first kernel **726**, a second kernel **728**, a third kernel **730**, and a fourth kernel **732**. The second processor **720** includes a fifth kernel **734**, a sixth kernel **736**, a seventh kernel **738**, and an eighth kernel **740**. The third processor **722** includes a ninth kernel **742**, a tenth kernel **744**, an eleventh kernel **746**, and a twelfth kernel **748**. The fourth processor **724** includes a thirteenth kernel **750**, a fourteenth kernel **752**, a fifteenth kernel **754**, and a sixteenth kernel **756**. Each of the kernels **726**, **728**, **730**, **732**, **734**, **736**, **738**, **740**, **742**, **744**, **746**, **748**, **750**, **752**, **754**, **756** may perform computations in parallel to generate output. The kernels **726**, **728**, **730**, **732**, **734**, **736**, **738**, **740**, **742**, **744**, **746**, **748**, **750**, **752**, **754**, **756** may be set to perform computations based on delimiter locations. For example, the kernels **726**, **728**, **730**, **732**, **734**, **736**, **738**, **740**, **742**, **744**, **746**, **748**, **750**, **752**, **754**, **756** may begin computations at data immediately following a delimiter. Fewer than all of the kernels **726**, **728**, **730**, **732**, **734**, **736**, **738**, **740**, **742**, **744**, **746**, **748**, **750**, **752**, **754**, **756** may be used in a particular execution cycle.

(78) Output of computations performed by the first processor **718** is placed into the first output buffer **758**. Output of computations performed by the second processor **720** is placed into the second output buffer **760**. Output of computations performed by the third processor **722** is placed into the third output buffer **762**. Output of computations performed by the fourth processor **724** is placed into the fourth output buffer **764**.

(79) Outputs in the output buffers **758**, **760**, **762**, **764** is gathered into the gathered output buffer **770** (e.g., by a downstream processor, such as a processor in the storage device compute module **508**) for output and/or additional processing.

(80) Thus, FIG. 7 illustrates an example in which a system performs carryover and forwarding operations based on first and last delimiter locations in order to keep host data units together for computation. Accordingly, parallel processing may be performed despite host data units being split

across multiple storage sources and processor input buffers.

(81) FIG. 8 is a diagram 800 showing examples of input buffer structures usable by the device 700 depicted in FIG. 7. The diagram 800 depicts a first buffer 802 following a copy from a storage source. The first buffer 802 corresponds to an input buffer for a first processing apparatus in a parallel processing stage. The first buffer 802 may correspond to the first input buffer 710. As shown, the first buffer 802 includes first extra buffer space 804 and second extra buffer space 806. The first extra buffer space 804 and the second extra buffer space 806 may have sizes based on a storage media data unit, based on a host data unit, or based on some other data unit. In some implementations, the first extra buffer space 804 corresponds to a third of the size of the first buffer 802 and the second extra buffer space 804 corresponds to a third of the size of the first buffer 802. Data from the storage source is copied in between the extra buffer spaces 804, 806. The data transferred into the first buffer 802 may be in storage media data units. An example storage media data unit 810 is shown. A size of the storage media data unit 810 may be different from a size of host data units utilized by a host application (e.g., an application data unit). An example complete application data unit 812 (e.g., a host data unit) is shown between a first application data delimiter 808 and a second application data delimiter 809. Because storage media data units and host data units may not be aligned, host data units may be split across storage sources. In the illustrated example, an incomplete host data unit is located before the first delimiter 808 and an incomplete host data unit is located after the second delimiter 809. The incomplete host data unit occurring before the first delimiter 808 may be completed by a carryover operation into the first extra buffer space 804 (e.g., as shown and described with reference to FIG. 7). Similarly, the host data unit occurring after the second delimiter 809 may be completed by a forwarding operation (e.g., as shown and described with reference to FIG. 7) into the second extra buffer space 806.

(82) The diagram 800 further depicts a second buffer 822 following a data copy from a storage source. The second buffer 822 corresponds to an input buffer for a processing apparatus in between first and last processing apparatuses in a parallel processing stage. The second buffer 822 may correspond to the second input buffer 712 or to the third input buffer 714. As shown, the second buffer 822 includes extra buffer space 826. The extra buffer space 826 may have a size based on a storage media data unit, based on a host data unit, or based on some other data unit. In some implementations the extra buffer space 826 is half of total space of the second buffer 822. In some implementations, the total space of the second buffer 822 is two thirds the size of the total space of the first buffer 802. Data from the storage source is copied into the second buffer 822 before the extra buffer space 826. The data transferred into the second buffer 822 may be in storage media data units. An example storage media data unit 830 is shown. A size of the storage media data unit 830 may be different from a size of host data units utilized by a host application (e.g., an application data unit). An example complete application data unit 832 (e.g., a host data unit) is shown between a first application data delimiter 828 and a second application data delimiter 829. Because storage media data units and host data units may not be aligned, host data units may be split across storage sources. In the illustrated example, an incomplete host data unit is located before the first delimiter 828 and an incomplete host data unit is located after the second delimiter 829. The incomplete host data unit occurring before the first delimiter 828 may be transferred by a forwarding operation into an extra buffer space of a previous input buffer (e.g., as shown and described with reference to FIG. 7). The data unit occurring after the second delimiter 829 may be completed by a forwarding operation from another input buffer into the extra buffer space 826 (e.g., as shown and described with reference to FIG. 7).

(83) The device 700 may have alternative configurations. For example, each of the input buffers 710, 712, 714, 716 may receive data from a different storage source than shown (e.g., a single media device, a group of storage channels, etc.). Additionally, the system may have more or fewer components (e.g., storage sources, compute modules, etc.) than shown. Further, processors may execute a different number of kernels than shown.

(84) The diagram **800** further depicts a third buffer **842** following a data copy from a storage source. The third buffer **842** corresponds to an input buffer for a last processing apparatus in a parallel processing stage. The third buffer **842** may correspond to the fourth input buffer **716**. A size of the third buffer **842** may be one third a total size of the first buffer **802** (e.g., because the third buffer lacks extra buffer space). Data from the storage source is copied into the third buffer **842**. The data transferred into the third buffer **842** may be in storage media data units. An example storage media data unit **850** is shown. A size of the storage media data unit **850** may be different from a size of host data units utilized by a host application (e.g., an application data unit). An example complete application data unit **852** (e.g., a host data unit) is shown between a first application data delimiter **848** and a second application data delimiter **849**. Because storage media data units and host data units may not be aligned, host data units may be split across storage sources. In the illustrated example, an incomplete host data unit is located before the first delimiter **848** and an incomplete host data unit is located after the second delimiter **849**. The incomplete host data unit occurring before the first delimiter **848** may be transferred by a forwarding operation into an extra buffer space of a previous input buffer (e.g., as shown and described with reference to FIG. 7). The data unit occurring after the second delimiter **849** transferred by a carryover operation into an extra buffer space of a first input buffer in a parallel processing stage (e.g., into the first extra buffer space **804**).

(85) As shown, first input buffers may maintain extra buffer space at a front end for carryover operations from a last input buffer. Further, input buffers before the last input buffer may maintain extra buffer space for forwarding operations. Accordingly, processors performing delimiter based parallel processing may reassemble complete host data units using carryover and forwarding operations, as described in FIG. 7.

(86) Referring to FIG. 9, a diagram showing a process **900** for performing delimiter based parallel processing is shown. The process **900** may be performed by a storage device or system, including the system **100** or the storage device **500**. The diagram depicts a first input buffer **902**, a second input buffer **904**, a third input buffer **906**, and a fourth input buffer **908**. The first input buffer **902** may correspond to the first input buffer **710**, the second input buffer **904** may correspond to the second input buffer **712**, the third input buffer **906** may correspond to the third input buffer **714**, and the fourth input buffer **908** may correspond to the fourth input buffer **716**.

(87) The process **900** includes a first DMA operation **920**, a first forwarding operation **922**, a processing operation **924**, a carryover operation **926**, a second DMA operation **928**, and a second forwarding operation **930**.

(88) In the first DMA operation **920**, first valid data **932** is copied into the first buffer **902** (e.g., by the first processor **718**) from a storage source (e.g., from the first storage source **702**). Further, second valid data **934** is copied into the second buffer **904** (e.g., by the second processor **720**) from a storage source (e.g., the second storage source **704**). Further, third valid data **936** is copied into the third buffer **906** (e.g., by the third processor **722**) from a storage source (e.g., the third storage source **706**). Further, fourth valid data **938** is copied into the fourth buffer **908** (e.g., by the fourth processor **724**) from a storage source (e.g., the fourth storage source **708**). As shown, the first buffer **902** has extra buffer space before (e.g., the first extra buffer space **804**) and after (e.g., the second extra buffer space **806**) the first valid data **932**. Additionally, the second buffer **904** and the third buffer **906** have extra buffer space after the second valid data **934** and the third valid data **936** respectively. During the DMA first DMA operation **920**, processors associated with the buffers **902**, **904**, **906**, **908** identify locations of delimiters within the valid data **932**, **934**, **936**, **938** (e.g., in the buffers **902**, **904**, **906**, **908**). In particular, the processors may identify locations of first and last delimiters within each of the buffers **902**, **904**, **906**, **908**.

(89) During the first forwarding operation **922**, a processor associated with the second buffer **904** identifies first data **942** before a first delimiter in the second valid data **934** and forwards the first data **942** to the first buffer **902** to form first modified valid data **940** when combined with the valid

data **932**.

(90) Further, a processor associated with the third buffer **906** identifies second data **946** before a first delimiter in the third valid data **936** and forwards the second data **946** to the second buffer **904**. Adding the second data **946** and removing the first data **942** from the second valid data **934** forms second modified valid data **944**.

(91) Further, a processor associated with the fourth buffer **908** identifies third data **950** before a first delimiter of the fourth valid data **952** and forwards the third data **950** to the third buffer **906**. Adding the third data **950** and removing the second data **946** from the third valid data **936** forms third modified valid data **948**. Removing the third data **950** from the fourth valid data **938** forms the fourth modified valid data **952**.

(92) In the first processing operation **924**, the processor associated with the first buffer performs computations based on the first modified valid data **940**. The processor associated with the second buffer **904** performs computations based on the second modified valid data **944**. The processor associated with the third buffer **906** performs computations based on the third modified valid data **948**. The processor associated with the fourth buffer **908** performs computations based on the fourth modified valid data **952** up to a final delimiter in the fourth modified valid data **952**.

(93) In the carryover operation **926**, the processor associated with the fourth buffer **908** identifies fourth data **954** following a final delimiter in the fourth modified valid data **952** and transfers (e.g., carries over) the fourth data **954** to the first buffer **902** (e.g., into extra space at the front of the first buffer **902**).

(94) In the second DMA operation **928**, the processor associated with the first buffer **902** copies data from a storage source into the first buffer to form fifth valid data **956** when combined with the fourth data **954**. The processor associated with the second buffer **904** copies data from a storage source into the second buffer **904** to form sixth valid data **958**. The processor associated with the third buffer **906** copies data from a storage source into the third buffer **906** to form seventh valid data **960**. The processor associated with the fourth buffer **908** copies data from a storage source into the fourth buffer to form eighth valid data **962**. The processors identify first and last delimiter locations in the buffers **902**, **904**, **906**, **908** during the second DMA operation **928**.

(95) In the second forwarding operation **930**, the processor associated with the second buffer **904** identifies fifth data **966** occurring before a first delimiter in the sixth valid data **958** and forwards the fifth data **966** to form fifth modified valid data **964**. The processor associated with the third buffer **906** identifies sixth data **970** occurring before a first delimiter in the seventh valid data **960** and forwards the sixth data **970** to the second buffer **904**. Adding the sixth data **970** to the sixth valid data **958** and subtracting the fifth data **966** forms sixth modified valid data **968**. The processor associated with the fourth buffer **908** may identify seventh data **974** before a first delimiter in the eighth valid data **962** and forward the seventh data **974** to the third buffer **906**. Adding the seventh data **974** and subtracting the sixth data **970** from the seventh valid data **960** may form seventh modified valid data **972**. Subtracting the seventh data **974** from the valid data **962** may form eighth modified valid data **976** in the fourth buffer **908**.

(96) The process **900** may continue on with additional processing, carryover, DMA, forwarding operations, etc. until a target amount of data is processed. Delimiter based execution together with the forwarding operations and the carryover operations may keep host data units together in one buffer despite how the host data units are split across storage sources. Accordingly, the process **900** may provide a technique for parallel processing in a system in which host data units are not aligned with a parallel processing pipeline.

(97) FIG. **10** is a diagram illustrating an example of a device **1000** performing delimiter based parallel computations in which compute modules operating in parallel pass data to one or more downstream processors based on delimiters. The device **1000** may correspond to the device **500** or to the system **100**.

(98) The device **1000** includes a first input buffer **1062**, a second input buffer **1064**, a third input

buffer **1066**, and a fourth input buffer **1068**. A first processor **1002** is associated with the first input buffer **1062**, a second processor **1004** is associated with the second input buffer **1064**, a third processor **1006** is associated with the third input buffer **1066**, and a fourth processor **1008** is associated with the fourth input buffer **1068**. The first processor **1002** is associated with a first output buffer **1070**, the second processor **1004** is associated with a second output buffer **1073**, the third processor **1006**, is associated with a third output buffer **1075**, and the fourth processor **1008** is associated with a fourth output buffer **1077**.

(99) The device **1000** further includes a first downstream input buffer **1011** associated with a first downstream processor **1010**. The first downstream processor **1010** is further associated with a first downstream output buffer **1072**. The device **1000** further includes a second downstream input buffer **1013** associated with a second downstream processor **1012**. The second downstream processor **1012** is further associated with a second downstream output buffer **1074**. The device **1000** further includes a third downstream input buffer **1015** associated with a third downstream processor **1014**. The third downstream processor **1014** is further associated with a third downstream output buffer **1076**. The device **1000** further includes fourth downstream input buffer **1017** associated with a fourth downstream processor **1016**. The fourth downstream processor **1016** is further associated with a fourth downstream output buffer **1078**.

(100) The first processor **1002**, first input buffer **1062**, and the first output buffer **1070** may correspond to a compute module, such as the compute module **600**. Similarly other combinations of associated processor, input buffer, and output buffer included in the device **1000** may correspond to compute modules, such as the compute module **600**.

(101) The downstream processors **1010**, **1012**, **1014**, **1016** are arranged downstream (e.g., closer to an output stage of a processing pipeline) from the processors **1002**, **1004**, **1006**, **1008**. In some examples, the first input buffer **1062**, the first processor **1002**, and the first output buffer **1070** correspond to the first channel compute module **518**; the second input buffer **1064**, the second processor **1004**, and the second output buffer **1073** correspond to the second channel compute module **520**; the third input buffer **1066**, the third processor **1006**, and the third output buffer **1075** correspond to the third channel compute module **526**; and the fourth input buffer **1068**, the fourth processor **1008**, and the fourth output buffer **1077** correspond to the fourth channel compute module **528**. The first downstream input buffer **1011**, the first downstream processor **1010**, and the first downstream output buffer **1072** may correspond to a downstream compute module, such as the storage device compute module **508**. Similarly, the second downstream input buffer **1013**, the second downstream processor **1012**, and the second downstream output buffer **1074** may correspond to the storage device compute module **508** (e.g., a device may have more than one storage device compute module). Further, the third downstream input buffer **1015**, the third downstream processor **1014**, and the third downstream output buffer **1076** may correspond to the storage device compute module **508**. Further, the fourth downstream input buffer **1017**, the fourth downstream processor **1016**, and the fourth downstream output buffer **1078** may correspond to the storage device compute module **508**.

(102) The device **1000** further includes a first storage source **1020**, a second storage source **1022**, a third storage source **1024**, a fourth storage source **1026**. The storage sources **1020**, **1022**, **1024**, **1026** may include storage media, storage media channels, storage media channel groups, etc. In some implementations, the first storage source **1020** corresponds to the first channel **534**, the second storage source **1022** corresponds to the second channel **536**, the third storage source **1024** corresponds to the third channel **538**, and the fourth storage source **1026** corresponds to the fourth channel **540**.

(103) The device **1000** further includes a gathered output buffer **1080**. The gathered output buffer **1080** may correspond to an output buffer of a compute module downstream from the processors illustrated.

(104) In operation, the first input buffer **1062** receives data from the first storage source **1020**. The

data may be placed in the first input buffer **1062** by a DMA operation performed by the first processor **1002**. The first processor **1002** may identify first and last delimiters within the data during the DMA operation. Similarly, the second input buffer **1064** receives data from the second storage source **1022**. The data may be placed in the second input buffer **1064** by a DMA operation performed by the second processor **1004**. The second processor **1004** may identify first and last delimiters within the data during the DMA operation. Further, the third input buffer **1066** receives data from the third storage source **1024**. The data may be placed in the third input buffer **1066** by a DMA operation performed by the third processor **1006**. The third processor **1006** may identify first and last delimiters within the data during the DMA operation. Further, the fourth input buffer **1068** receives data from the fourth storage source **1026**. The data may be placed in the fourth input buffer **1068** by a DMA operation performed by the fourth processor **1008**. The fourth processor **1008** may identify first and last delimiters within the data during the DMA operation.

(105) The first processor **1002** sends data before a first delimiter in the first input buffer **1062** to the fourth downstream input buffer **1017**. Further, the processor **1002** sends data after a last delimiter in the first input buffer **1062** to the first downstream input buffer **1011**. The second processor **1004** sends data before a first delimiter in the second input buffer **1064** to the first downstream input buffer **1011**. Accordingly, a host data unit split across the first input buffer **1062** and the second input buffer **1064** is put back together in the first downstream input buffer **1011**. The second processor **1012** further sends data occurring after a last delimiter in the second input buffer **1064** to the second downstream input buffer **1013**. The third processor **1006** sends data occurring before a first delimiter in the third input buffer **1066** to the second downstream input buffer **1013**.

Accordingly, a host data unit split across the second input buffer **1064** and the third input buffer **1066** is put back together in the second downstream input buffer **1013**. The third processor **1006** sends data occurring after a last delimiter in the third input buffer **1066** to the third downstream input buffer **1015**. The fourth processor **1008** sends data occurring before a first delimiter in the fourth input buffer **1068** to the third downstream input buffer **1015**. Accordingly, a host data unit split across the third input buffer **1066** and the fourth input buffer **1068** is put back together in the third downstream input buffer **1015**. The fourth processor **1008** sends data occurring after a last delimiter in the fourth input buffer **1068** to the fourth downstream input buffer **1017**. Accordingly, a host data unit split across the fourth input buffer **1068** and the first input buffer **1062** is put back together in the fourth downstream input buffer **1017**.

(106) Each of the processors **1002**, **1004**, **1006**, **1008** in the device **1000** includes four kernels. Further each of the downstream processors **1010**, **1012**, **1016** includes a kernel. A different number of kernels may be included than shown. Each kernel may include hardware and/or executable software configured to perform a computation on data (e.g., a search function, a filter function, a mathematical operation, etc.) to generate output. The first processor **1002** includes a first kernel **1030**, a second kernel **1032**, a third kernel **1034**, and a fourth kernel **1036**. The second processor **1004** includes a fifth kernel **1038**, a sixth kernel **1040**, a seventh kernel **1042**, and an eighth kernel **1044**. The third processor **1006** includes a ninth kernel **1046**, a tenth kernel **1048**, an eleventh kernel **1050**, and a twelfth kernel **1052**. The fourth processor **1008** includes a thirteenth kernel **1054**, a fourteenth kernel **1056**, a fifteenth kernel **1058**, and a sixteenth kernel **1060**. Each of the kernels **1030**, **1032**, **1034**, **1036**, **1038**, **1040**, **1042**, **1044**, **1046**, **1048**, **1050**, **1052**, **1054**, **1056**, **1058**, **1060** may perform computations in parallel to generate output.

(107) The first downstream processor **1010** includes a first downstream kernel **1082**, the second downstream processor **1012** includes a second downstream kernel **1084**, the third downstream processor **1014** includes a third downstream kernel **1086**, and the fourth processor **1016** includes a fourth downstream kernel **1088**. Each of the downstream kernels **1082**, **1084**, **1086**, **1088** may operate in parallel to generate output. In some examples, the kernels **1030**, **1032**, **1034**, **1038**, **1040**, **1042**, **1044**, **1046**, **1048**, **1050**, **1052**, **1054**, **1056**, **1058**, **1060** may operate in parallel with the downstream kernels **1082**, **1084**, **1086**, **1088**.

(108) The kernels **1030, 1032, 1034, 1036, 1038, 1040, 1042, 1044, 1046, 1048, 1050, 1052, 1054, 1056, 1058, 1060** and the downstream kernels **1082, 1084, 1086, 1088** may be set to perform computations based on delimiter locations. For example, the kernels **1030, 1032, 1034, 1036, 1038, 1040, 1042, 1044, 1046, 1048, 1050, 1052, 1054, 1056, 1058, 1060** may begin computations at data immediately following a delimiter. The downstream kernels **1082, 1084, 1086, 1088** may operate in a similar fashion. Fewer than all of the kernels **1030, 1032, 1034, 1036, 1038, 1040, 1042, 1044, 1046, 1048, 1050, 1052, 1054, 1056, 1058, 1060** may be used in some examples.

(109) Output of computations performed by the first processor **1002** is placed into the first output buffer **1070**. Output of computations performed by the second processor **1004** is placed into the second output buffer **1073**. Output of computations performed by the third processor **1006** is placed into the third output buffer **1075**. Output of computations performed by the fourth processor **1008** is placed into the fourth output buffer **1077**. Output from the first downstream processor **1010** is placed into the first downstream output buffer **1072**. Output from the second downstream processor **1012** is placed into the second downstream output buffer **1074**. Output from the third downstream processor **1014** is placed into the third downstream output buffer **1076**. Output from the fourth downstream processor **1016** into the fourth downstream output buffer **1078**.

(110) Outputs in the output buffers **1070, 1073, 1075, 1077** and the downstream output buffers **1072, 1074, 1076, 1078** is gathered into the gathered output buffer **1080** (e.g., by a downstream processor, such as a processor in the storage device compute module **508**) for output and/or additional processing.

(111) Because incomplete host data units located at ends of input buffers are put back together in downstream input buffers based on delimiter locations, the device **1000** provides efficient parallel processing of host data that is stored in a fashion that is unaligned with a parallel processing pipeline.

(112) FIG. **11** is a diagram illustrating an example of a device **1100** performing delimiter based parallel computations in which compute modules operating in parallel directly access input buffers of other compute modules based on delimiter locations. The device **1100** is a computational storage device and may correspond to the device **500** of FIG. **5**. The device **1100** includes a first processor **1110** associated with a first input buffer **1120**, a second processor **1112** associated with a second input buffer **1122**, a third processor **1114** associated with a third input buffer **1124**, and a fourth processor **1116** associated with a fourth input buffer **1126**. In some examples, the first processor **1110** and the first input buffer **1120** correspond to the first channel compute module **518**, the second processor **1112** and the second input buffer **1122** correspond to the second channel compute module **520**, the third processor **1114** and the third input buffer **1124** correspond to the third channel compute module **526**, and the fourth processor **1116** and the fourth input buffer **1126** correspond to the fourth channel compute module **528**. In other examples, the processors **1110, 1112, 1114, 1116** and input buffers **1120, 1122, 1124, 1126** correspond to different compute modules in the device **500**.

(113) In the device **1100**, the input buffers **1120, 1122, 1124, 1126** are ranges within a common buffer space. Data from a first storage source **1102** is copied to the first input buffer **1120** (e.g., by a DMA operation initiated by the first processor **1110**). Data from a second storage source **1104** is copied to the second input buffer **1122** from a second storage source, data from a third storage source **1106** is copied to a third input buffer **1124**, and data from a fourth storage source **1108** is copied into the fourth input buffer **1126**. The first processor **1110** identifies first and last delimiter locations within the first input buffer **1120**, the second processor **1112** identifies first and last delimiter locations within the second input buffer **1122**, the third processor **1114** identifies first and last delimiter locations within the third input buffer **1124**, and the fourth processor **1116** identifies first and last delimiter locations within the fourth input buffer **1126**.

(114) In operation, the first processor **1110** begins processing on data after a first delimiter in the first input buffer **1120** and continues processing into the second input buffer **1122** until a first

delimiter in the second input buffer **1122** is reached. Similarly, the second processor **1112** begins processing data occurring after the first delimiter in the second input buffer **1122** and continues processing into the third input buffer **1124** until a first delimiter in the third input buffer is reached. Similarly, the third processor **1114** begins processing data occurring after the first delimiter in the third input buffer **1124** and continues processing until reaching a first delimiter in the fourth input buffer **1126**. Similarly, the fourth input buffer **1126** begins processing data occurring after the first delimiter in the fourth input buffer and upon reaching an end of the fourth input buffer **1126**, begins processing at a beginning of the first input buffer **1120** and continues until reaching a first delimiter in the first input buffer **1120**.

(115) The processors **1110**, **1112**, **1114**, **1116** may operate in parallel. Further, because the processors may read past boundaries of their associated input buffers, a host data unit may be processed by a single processor rather than being split across processors. Accordingly, the device **1100** may provide parallel processing of data in systems in which host data units are not aligned with a parallel processing pipeline. It should be noted that input buffers of the device **1100** may not include extra buffer space as in some other implementations described herein.

(116) FIG. **12** is a flowchart of a method **1200** for delimiter based parallel processing. The method **1200** may be performed by the system **100**, by the device **500**, by the device **700**, by the device **1000**, or by the device **1100**.

(117) The method **1200** includes receiving, at a first buffer, first data from a first storage source, at **1202**. For example, the system first buffer **106** may receive the first data, including the first delimiter **113** and the first portion **114**, from the first storage source **110**. As another example, the first buffer **106** may receive the first data from the first storage source **110** via the third processor **302** and the third buffer **304**. As another example, the first input buffer **710** may receive data from the first storage source **702**. As another example, the first downstream input buffer **1011** may receive data from the first storage source **1020** via the first input buffer **1062** and the first processor **1002**. As another example, the first input buffer **1120** may receive data from the first storage source **1102**.

(118) The method **1200** further includes receiving, at a second buffer, second data from a second storage source, the second data including a first portion and a second portion separated by a delimiter, at **1204**. For example, the second buffer **108** may receive the second data, including the second portion **116**, the second delimiter **118**, and the third delimiter **121**, from the second storage source **112**. As another example, the second input buffer **712** may receive data from the second storage source **704**. As another example, the second input buffer **1064** may receive data from the second storage source **1022**. As another example, the second input buffer **1122** may receive data from the second storage source **1104**.

(119) The method **1200** further includes performing, at a second processor associated with the second buffer, a first operation on the second portion of the second data, at **1206**. For example, the second processor **104** may perform a computation on the third portion **120** in response to the third portion falling between a first and last delimiter in the second buffer **108** (e.g., the second delimiter **118** and the third delimiter **121**). As another example, the second processor **720** may perform computations on data in the second input buffer **712** that falls between a first and last delimiter within the second input buffer. As another example, the second processor **1004** may perform computations based on data in the second input buffer **1064** between a first and last delimiter within the second input buffer **1064**. As another example, the second processor **1112** may perform computations based on data in the second input buffer **1122** falling between a first and last delimiter within the second input buffer **1122**.

(120) The method **1200** further includes performing, at a first processor associated with the first buffer, a second operation on the first data and the first portion of the second data, at **1208**. For example, the second processor **104** may copy the second portion **116** into the first buffer **106** in response to the second portion **116** being located before a first delimiter (e.g., the second delimiter

118) in the second buffer **108**. The first processor **102** may perform computations based on the first portion **114** and the second portion **116**. As another example, the first processor **102** may perform computations beginning after a first delimiter in the first buffer **106** (e.g., the first delimiter **113**) and continuing into the second buffer **108** until reaching a first delimiter within the second buffer **108** (e.g., the second delimiter **118**). As another example, the second processor **720** may copy data located before a first delimiter in the second input buffer **712** into the first input buffer **710** and the first processor **718** may perform computations based on data in the first input buffer **710**. In another example, the second processor **1004** copies data before a first delimiter in the second buffer **1064** to the first downstream input buffer **1011** and the first downstream processor **1010** performs computations based on data in the first downstream input buffer **1011**. In another example, the first processor **1110** performs computations on data starting after a first delimiter in the first input buffer **1120** and continues into the second input buffer **1122** until reaching a first delimiter within the second input buffer **1122**.

(121) The method **1200** may be used to perform parallel processing in systems in which host data units are not aligned with storage data units and/or a parallel processing pipeline.

(122) In some examples, X corresponds to Y based on X matching Y. For example, a first ID may be determined to correspond to a second ID that matches (e.g., has a same value as) the first ID. In other examples, X correspond to Y based on X being associated with (e.g., linked to) Y. For example, X may be associated to Y by a mapping data structure.

(123) Certain embodiments may be implemented in one or a combination of hardware, firmware, and software. Other embodiments may also be implemented as instructions stored on a computer-readable storage device, which may be read and executed by at least one processor to perform the operations described herein. A computer-readable storage device may include any non-transitory memory mechanism for storing information in a form readable by a machine (e.g., a computer). For example, a computer-readable storage device may include read-only memory (ROM), random-access memory (RAM), magnetic disk storage media, optical storage media, flash-memory devices, and other storage devices and media.

(124) As used within this document, the term “communicate” is intended to include transmitting, or receiving, or both transmitting and receiving. This may be particularly useful in claims when describing the organization of data that is being transmitted by one device and received by another, but only the functionality of one of those devices is required to infringe the claim. Similarly, the bidirectional exchange of data between two devices (both devices transmit and receive during the exchange) may be described as ‘communicating’, when only the functionality of one of those devices is being claimed. The term “communicating” as used herein with respect to a wireless communication signal includes transmitting the wireless communication signal and/or receiving the wireless communication signal. For example, a wireless communication unit, which is capable of communicating a wireless communication signal, may include a wireless transmitter to transmit the wireless communication signal to at least one other wireless communication unit, and/or a wireless communication receiver to receive the wireless communication signal from at least one other wireless communication unit.

(125) Some embodiments may be used in conjunction with various devices and systems, for example, a Personal Computer (PC), a desktop computer, a mobile computer, a laptop computer, a notebook computer, a tablet computer, a server computer, a handheld computer, a handheld device, a Personal Digital Assistant (PDA) device, a handheld PDA device, an on-board device, an off-board device, a hybrid device, a vehicular device, a non-vehicular device, a mobile or portable device, a consumer device, a non-mobile or non-portable device, a wireless communication station, a wireless communication device, a wireless Access Point (AP), a wired or wireless router, a wired or wireless modem, a video device, an audio device, an audio-video (A/V) device, a wired or wireless network, a wireless area network, a Wireless Video Area Network (WVAN), a Local Area Network (LAN), a Wireless LAN (WLAN), a Personal Area Network (PAN), a Wireless PAN

(WPAN), and the like.

(126) Some embodiments may be used in conjunction with one way and/or two-way radio communication systems, cellular radio-telephone communication systems, a mobile phone, a cellular telephone, a wireless telephone, a Personal Communication Systems (PCS) device, a PDA device which incorporates a wireless communication device, a mobile or portable Global Positioning System (GPS) device, a device which incorporates a GPS receiver or transceiver or chip, a device which incorporates an radio frequency identification (RFID) element or chip, a Multiple Input Multiple Output (MIMO) transceiver or device, a Single Input Multiple Output (SIMO) transceiver or device, a Multiple Input Single Output (MISO) transceiver or device, a device having one or more internal antennas and/or external antennas, Digital Video Broadcast (DVB) devices or systems, multi-standard radio devices or systems, a wired or wireless handheld device, e.g., a Smartphone, a Wireless Application Protocol (WAP) device, or the like.

(127) Some embodiments may be used in conjunction with one or more types of wireless communication signals and/or systems following one or more wireless communication protocols, for example, Radio Frequency (RF), Infrared (IR), Frequency-Division Multiplexing (FDM), Orthogonal FDM (OFDM), Time-Division Multiplexing (TDM), Time-Division Multiple Access (TDMA), Extended TDMA (E-TDMA), General Packet Radio Service (GPRS), extended GPRS, Code-Division Multiple Access (CDMA), Wideband CDMA (WCDMA), CDMA 2000, single-carrier CDMA, multi-carrier CDMA, Multi-Carrier Modulation (MDM), Discrete Multi-Tone (DMT), Bluetooth™, Global Positioning System (GPS), Wi-Fi, Wi-Max, ZigBee™, Ultra-Wideband (UWB), Global System for Mobile communication (GSM), 2G, 2.5G, 3G, 3.5G, 4G, Fifth Generation (5G) mobile networks, 3GPP, Long Term Evolution (LTE), LTE advanced, Enhanced Data rates for GSM Evolution (EDGE), or the like. Other embodiments may be used in various other devices, systems, and/or networks.

(128) Although an example processing system has been described above, embodiments of the subject matter and the functional operations described herein can be implemented in other types of digital electronic circuitry, or in computer software, firmware, or hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them.

(129) Embodiments of the subject matter and the operations described herein can be implemented in digital electronic circuitry, or in computer software, firmware, or hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them. Embodiments of the subject matter described herein can be implemented as one or more computer programs, i.e., one or more components of computer program instructions, encoded on computer storage medium for execution by, or to control the operation of, information/data processing apparatus. Alternatively, or in addition, the program instructions can be encoded on an artificially-generated propagated signal, for example a machine-generated electrical, optical, or electromagnetic signal, which is generated to encode information/data for transmission to suitable receiver apparatus for execution by an information/data processing apparatus. A computer storage medium can be, or be included in, a computer-readable storage device, a computer-readable storage substrate, a random or serial access memory array or device, or a combination of one or more of them. Moreover, while a computer storage medium is not a propagated signal, a computer storage medium can be a source or destination of computer program instructions encoded in an artificially-generated propagated signal. The computer storage medium can also be, or be included in, one or more separate physical components or media (for example multiple CDs, disks, or other storage devices).

(130) The operations described herein can be implemented as operations performed by an information/data processing apparatus on information/data stored on one or more computer-readable storage devices or received from other sources.

(131) The term “data processing apparatus” encompasses all kinds of apparatus, devices, and

machines for processing data, including by way of example a programmable processor, a computer, a system on a chip, or multiple ones, or combinations, of the foregoing. The apparatus can include special purpose logic circuitry, for example an FPGA (field programmable gate array) or an ASIC (application-specific integrated circuit). The apparatus can also include, in addition to hardware, code that creates an execution environment for the computer program in question, for example code that constitutes processor firmware, a protocol stack, a database management system, an operating system, a cross-platform runtime environment, a virtual machine, or a combination of one or more of them. The apparatus and execution environment can realize various different computing model infrastructures, such as web services, distributed computing and grid computing infrastructures.

(132) A computer program (also known as a program, software, software application, script, or code) can be written in any form of programming language, including compiled or interpreted languages, declarative or procedural languages, and it can be deployed in any form, including as a stand-alone program or as a component, component, subroutine, object, or other unit suitable for use in a computing environment. A computer program may, but need not, correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or information/data (for example one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (for example files that store one or more components, sub-programs, or portions of code). A computer program can be deployed to be executed on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

(133) The processes and logic flows described herein can be performed by one or more programmable processors executing one or more computer programs to perform actions by operating on input information/data and generating output. Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and information/data from a read-only memory or a random access memory or both. Elements of a computer are a processor for performing actions in accordance with instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive information/data from or transfer information/data to, or both, one or more mass storage devices for storing data, for example magnetic, magneto-optical disks, or optical disks. However, a computer need not have such devices. Devices suitable for storing computer program instructions and information/data include all forms of non-volatile memory, media and memory devices, including by way of example semiconductor memory devices, for example EPROM, EEPROM, and flash memory devices; magnetic disks, for example internal hard disks or removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(134) To provide for interaction with a user, embodiments of the subject matter described herein can be implemented on a computer having a display device, for example a CRT (cathode ray tube) or LCD (liquid crystal display) monitor, for displaying information/data to the user and a keyboard and a pointing device, for example a mouse or a trackball, by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, for example visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's client device in response to requests received from the web browser.

(135) Embodiments of the subject matter described herein can be implemented in a computing

system that includes a back-end component, for example as an information/data server, or that includes a middleware component, for example an application server, or that includes a front-end component, for example a client computer having a graphical user interface or a web browser through which a user can interact with an embodiment of the subject matter described herein, or any combination of one or more such back-end, middleware, or front-end components. The components of the system can be interconnected by any form or medium of digital information/data communication, for example a communication network. Examples of communication networks include a local area network (“LAN”) and a wide area network (“WAN”), an inter-network (for example the Internet), and peer-to-peer networks (for example ad hoc peer-to-peer networks).

(136) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship to each other. In some embodiments, a server transmits information/data (for example an HTML page) to a client device (for example for purposes of displaying information/data to and receiving user input from a user interacting with the client device). Information/data generated at the client device (for example a result of the user interaction) can be received from the client device at the server.

(137) While this specification contains many specific embodiment details, these should not be construed as limitations on the scope of any embodiment or of what may be claimed, but rather as descriptions of features specific to particular embodiments. Certain features that are described herein in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable sub-combination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a sub-combination or variation of a sub-combination.

(138) Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

(139) Thus, particular embodiments of the subject matter have been described. Other embodiments are within the scope of the following claims. In some cases, the actions recited in the claims can be performed in a different order and still achieve desirable results. In addition, the processes depicted in the accompanying figures do not necessarily require the particular order shown, or sequential order, to achieve desirable results. In certain embodiments, multitasking and parallel processing may be advantageous.

(140) Many modifications and other embodiments of the disclosure set forth herein will come to mind to one skilled in the art to which these embodiments pertain having the benefit of the teachings presented in the foregoing descriptions and the associated drawings. Therefore, it is to be understood that the embodiments are not to be limited to the specific embodiments disclosed and that modifications and other embodiments are intended to be included within the scope of the appended claims. Although specific terms are employed herein, they are used in a generic and descriptive sense only and not for purposes of limitation.

Claims

1. A storage device comprising: a first storage source storing first data; a second storage source storing second data, the second data including a first portion and a second portion separated by a delimiter; a first buffer configured to receive the first data; a second buffer configured to receive the second data and store the first portion, the second portion, and the delimiter; a first processor associated with the first buffer; and a second processor associated with the second buffer, wherein the second processor is configured to perform a first operation on the second portion of the second data, wherein the first processor is configured to: receive a request from a third processor of a host computing device to perform a computation on a data unit, the third processor being different from the first processor and the second processor; identify, based on the request and further based on the delimiter, the first data and the first portion of the second data as comprising the data unit; perform a second operation on the first data and the first portion of the second data based on the delimiter and generate an output, wherein the second operation includes a first computation based on the first data and a second computation based on the first portion of the second data, wherein the first processor is configured to access the first portion of the second data for performing the second operation based on a location of the delimiter; and return the output to the third processor in response to the request.
2. The storage device of claim 1, wherein the second processor is configured to copy the first portion of the second data from the second buffer to the first buffer.
3. The storage device of claim 1, wherein the delimiter corresponds to a first delimiter in the second buffer, and wherein the second processor is configured to copy the first portion of the second data to the first buffer in response to the first portion being located before the first delimiter within the second buffer.
4. The storage device of claim 3, wherein the first buffer includes a first extra buffer space in front of the first data in the first buffer and a second extra buffer space following the first data in the first buffer, and wherein the second processor is configured to copy the first portion to the second extra buffer space.
5. The storage device of claim 4, further comprising: a third buffer storing third data; and a third processor associated with the third buffer, wherein the third processor is configured to copy the third data to the first extra buffer space in response to the third data being located after a final delimiter in the third buffer.
6. The storage device of claim 3, wherein the first data follows a final delimiter in the first buffer, and wherein the first processor is configured to perform the second operation in response to the first buffer receiving the first portion of the second data.
7. The storage device of claim 1, wherein the storage device includes a buffer space that includes the first buffer and the second buffer, and wherein the first processor is configured to stop processing in the second buffer based on a location of the delimiter, and wherein the second processor is configured to begin processing in the second buffer based on the location of the delimiter.
8. The storage device of claim 1, wherein the first storage source comprises a storage channel, a storage media device, or a group of storage channels.
9. The storage device of claim 1, wherein the first storage source comprises a NAND flash channel, a NAND flash chip, or a group of NAND flash channels.
10. The storage device of claim 1, further comprising: a third buffer configured to receive the first data from the first storage source; and a third processor associated with the third buffer configured to copy the first data from the third buffer to the first buffer in response to the first data being located after a final delimiter in the third buffer.
11. A method comprising: receiving, at a first buffer, first data from a first storage source;

receiving, at a second buffer, second data from a second storage source, the second data including a first portion and a second portion separated by a delimiter; storing the first portion, the second portion, and the delimiter in the second buffer; performing, at a second processor associated with the second buffer, a first operation on the second portion of the second data; receiving, by a first processor associated with the first buffer, a request from a third processor of a host computing device to perform a computation on a data unit, the third processor being different from the first processor and the second processor; identifying, by the first processor, based on the request and further based on the delimiter, the first data and the first portion of the second data as comprising the data unit; performing, by the first processor, a second operation on the first data and the first portion of the second data based on the delimiter and generating an output, wherein the second operation includes a first computation based on the first data and a second computation based on the first portion of the second data, wherein the first processor is configured to access the first portion of the second data for performing the second operation based on a location of the delimiter; and returning the output to the third processor in response to the request.

12. The method of claim 11, further comprising copying the first portion of the second data from the second buffer to the first buffer.

13. The method of claim 11, wherein the delimiter corresponds to a first delimiter in the second buffer, the method further comprising copying the first portion of the second data to the first buffer in response to the first portion being located before the first delimiter within the second buffer.

14. The method of claim 13, wherein the first buffer includes a first extra buffer space in front of the first data in the first buffer and a second extra buffer space following the first data in the first buffer, and wherein copying the first portion of the second data to the first buffer includes copying the first portion to the second extra buffer space.

15. The method of claim 14, further comprising: storing third data in a third buffer associated with a third processor; and copying the third data to the first extra buffer space in response to the third data being located after a final delimiter in the third buffer.

16. The method of claim 11, wherein the first buffer and the second buffer are included in a common buffer space, and wherein performing the second operation includes: the first processor performing operations starting in the first buffer and stopping in the second buffer based on a location of the delimiter.

17. The method of claim 11, wherein the first storage source comprises a storage channel, a storage media device, or a group of storage channels.

18. The method of claim 11, wherein the first storage source comprises a NAND flash channel, a NAND flash chip, or a group of NAND flash channels.

19. The method of claim 11, further comprising: receive the first data from the first storage source at a third buffer; and copying the first data from the third buffer to the first buffer in response to the first data being located after a final delimiter in the third buffer.

20. A storage device comprising: a first storage channel including a first media device storing first data; a second storage channel including a second media device storing second data, the second data including a first portion and a second portion separated by a delimiter; a first compute module associated with the first storage channel and including a first processor and a first input buffer; and a second compute module associated with the second storage channel and including a second processor and a second input buffer, wherein the second input buffer is configured to store the first portion, the second portion, and the delimiter, wherein the second processor is configured to perform a first operation on the second portion of the second data, wherein the first processor is configured to: receive a request from a third processor of a host computing device to perform a computation on a data unit, the third processor being different from the first processor and the second processor; identify, based on the request and further based on the delimiter, the first data and the first portion of the second data as comprising the data unit; perform a second operation on the first data and the first portion of the second data based on the delimiter and generate an output,

wherein the second operation includes a first computation based on the first data and a second computation based on the first portion of the second data, wherein the first processor is configured to access the first portion of the second data for performing the second operation based on a location of the delimiter; and return the output to the third processor in response to the request.
