

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12411670
B2
September 09, 2025
Rister; Blaine et al.

Analysis assistant for determining execution inefficiencies in dataflow programs

Abstract

In a method, in response to an interface a computer-implemented analysis assistant initiates a presentation of inefficiency results, determined an efficiency analyzer based on a mapping of a dataflow program to execute on hardware of a computing system. The assistant receives an inefficiency included among the inefficiency results and composes formatted inefficiency results comprising a presentation format of the inefficiency to assist a developer of the dataflow program to interpret the inefficiency. The analysis assistant outputs the formatted inefficiency results to an interface, which can comprise an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency. In implementations the presentation can comprise an interactive presentation with a developer of the dataflow program. A computer program product and a computing system can implement the method. The computing system can execute the assistant, and can include the interfaces.

Inventors: Rister; Blaine (San Francisco, CA), Li; Qingjian (Berkeley, CA), Yang; Bowen (Palo Alto, CA), Wang; Junjue (Newark, CA), Liu; Chen (Palo Alto, CA), Chen; Zhuo (Palo Alto, CA), Sujeeth; Arvind (Palo Alto, CA), Jairath; Sumti (Palo Alto, CA)

Applicant: SambaNova Systems, Inc. (Palo Alto, CA)

Family ID: 88194358

Assignee: SambaNova Systems, Inc. (Palo Alto, CA)

Appl. No.: 18/387912

Filed: November 08, 2023

Prior Publication Data

Document Identifier
US 20240078098 A1

Publication Date
Mar. 07, 2024

Related U.S. Application Data

continuation parent-doc US 18129718 20230331 PENDING child-doc US 18387912
us-provisional-application US 63331696 20220415
us-provisional-application US 63331116 20220414
us-provisional-application US 63330740 20220413
us-provisional-application US 63330730 20220413
us-provisional-application US 63327313 20220404
us-provisional-application US 63326762 20220401
us-provisional-application US 63326206 20220331

Publication Classification

Int. Cl.: G06F9/44 (20180101); G06F8/41 (20180101)

U.S. Cl.:

CPC G06F8/433 (20130101);

Field of Classification Search

CPC: G06F (8/433)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11080227	12/2020	Koeplinger et al.	N/A	N/A
11195080	12/2020	Nama et al.	N/A	N/A
11227207	12/2021	Nama et al.	N/A	N/A
11410027	12/2021	Chen et al.	N/A	N/A
11500871	12/2021	Kulkarni	N/A	G06N 20/00
2014/0237459	12/2013	Von Platen	717/155	G06F 8/452
2019/0042217	12/2018	Glossop	N/A	G06F 8/443
2021/0365248	12/2020	ChoFleming, Jr.	N/A	G06F 8/433
2022/0012028	12/2021	Yount	N/A	G06F 8/443

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2010142987	12/2009	WO	N/A

OTHER PUBLICATIONS

Koeplinger et al., Spatial: A Language and Compiler for Application Accelerators, PLDI '18, Jun. 18-22, 2018, Association for Computing Machinery, 16 pages. cited by applicant
List of related cases—SBNV1114USC01, Jun. 6, 2024, 2 pages. cited by applicant
M. Emani et al., Accelerating Scientific Applications With Sambanova Reconfigurable Dataflow Architecture, in Computing in Science & Engineering, vol. 23, No. 2, pp. 114-119, Mar. 26, 2021, [doi: 10.1109/MCSE.2021.3057203]. cited by applicant
Podobas et al, A Survey on Coarse-Grained Reconfigurable Architectures From a Performance

Perspective, IEEEAccess, vol. 2020.3012084, Jul. 27, 2020, 25 pages. cited by applicant
Prabhakar et al., Plasticine: A Reconfigurable Architecture for Parallel Patterns, ISCA, Jun. 24-28, 2017, 14 pages. cited by applicant
U.S. Appl. No. 16/239,252 Final Office Action, dated Jan. 8, 2020, 13 pages. cited by applicant
U.S. Appl. No. 16/239,252—Notice of Allowance dated Feb. 12, 2020, 10 pages. cited by applicant
U.S. Appl. No. 16/239,252—Notice of Allowance dated May 14, 2020, 15 pages. cited by applicant
U.S. Appl. No. 16/239,252—Office Action dated Aug. 7, 2019, 8 pages. cited by applicant
U.S. Appl. No. 16/239,252—Response to Final Office Action dated Jan. 8, 2020, filed Jan. 24, 2020, 14 pages. cited by applicant
U.S. Appl. No. 16/239,252—Response to Office Action dated Aug. 7, 2019, filed Sep. 26, 2019, 6 pages. cited by applicant
U.S. Appl. No. 17/216,651 Non-Final Rejection, dated Jul. 13, 2021, 12 pages. cited by applicant
U.S. Appl. No. 17/216,651 Notice of Allowance, dated Aug. 5, 2021, 14 pages. cited by applicant
U.S. Appl. No. 17/216,651 Response to First Office Action, dated Jul. 13, 2021, filed Jul. 23, 2021, 14 pages. cited by applicant
U.S. Appl. No. 17/216,652 Non-Final Rejection, dated Aug. 2, 2021, 24 pages. cited by applicant

Primary Examiner: Nahar; Qamrun

Attorney, Agent or Firm: Flagship Patents

Background/Summary

CROSS-REFERENCE AND INCORPORATIONS (1) This application claims the benefit of U.S. Non-Provisional Patent Application No. 18/129,718 filed Mar. 31, 2023 which further claims benefit of U.S. Provisional Patent Application No. 63/327,313 filed Apr. 4, 2022, which is incorporated by reference herein in its entirety. (2) This application further claims the benefit of U.S. Provisional Patent Application No. 63/330,730 filed Apr. 13, 2022, which is incorporated by reference herein in its entirety. (3) This application further claims the benefit of U.S. Provisional Patent Application No. 63/330,740 filed Apr. 13, 2022, which is incorporated by reference herein in its entirety. (4) This application further claims the benefit of U.S. Provisional Patent Application No. 63/326,206 filed Mar. 31, 2022, which is incorporated by reference herein in its entirety. (5) This application further claims the benefit of U.S. Provisional Patent Application No. 63/326,762 filed Apr. 1, 2022, which is incorporated by reference herein in its entirety. (6) This application further claims the benefit of U.S. Provisional Patent Application No. 63/331,116 filed Apr. 14, 2022, which is incorporated by reference herein in its entirety. (7) This application further claims the benefit of U.S. Provisional Patent Application No. 63/331,696 filed Apr. 15, 2022, which is incorporated by reference herein in its entirety. (8) The following are incorporated by reference for all purposes as if fully set forth herein: Prabhakar et al., “Plasticine: A Reconfigurable Architecture for Parallel Patterns,” ISCA '17, Jun. 24-28, 2017, Toronto, ON, Canada; Koeplinger et al., “Spatial: A Language And Compiler For Application Accelerators,” Proceedings Of The 39th ACM SIGPLAN Conference On Programming Language Design And Embodiment (PLDI), Proceedings of the 43rd International Symposium on Computer Architecture, 2018; U.S. Nonprovisional patent application Ser. No. 16/239,252, filed Jan. 3, 2019, titled “VIRTUALIZATION OF A RECONFIGURABLE DATA PROCESSOR”; U.S. Nonprovisional patent application Ser. No. 16/536,192, filed Aug. 8, 2019, entitled “COMPILER FLOW LOGIC FOR RECONFIGURABLE ARCHITECTURES”; U.S. Nonprovisional patent application Ser. No. 16/572,527, filed Sep. 16, 2019, entitled “PERFORMANCE ESTIMATION-BASED RESOURCE ALLOCATION FOR

RECONFIGURABLE ARCHITECTURES”; U.S. Nonprovisional patent application Ser. No. 17/216,651, filed Mar. 29, 2021, entitled “LOSSLESS TILING IN CONVOLUTION NETWORKS—TILING CONFIGURATION”; and, U.S. Nonprovisional patent application Ser. No. 17/216,652, filed Mar. 29, 2021, entitled “LOSSLESS TILING IN CONVOLUTION NETWORKS—SECTION BOUNDARIES”. FIELD OF THE TECHNOLOGY (9) The technology disclosed relates to dataflow computing, such as neural networks in machine learning and artificial intelligence computing systems. In particular, the technology disclosed relates to compilers for data parallel and dataflow computing systems, and computing systems using reconfigurable processors, such as coarse-grain reconfigurable processors (CGRPs) to execute dataflow computing applications.

BACKGROUND

(1) Advanced computing applications, such as neural networks, machine learning, and artificial intelligence applications, can be executed by dataflow and/or data parallel computing systems. The present disclosure (hereinafter, “the disclosure”) relates to such applications and computing systems for executing such applications. In particular, the disclosure relates to program compilers of computing systems, and to compiler optimization of allocation of hardware resources of computing systems to execute functions of such applications.

(2) Computing systems can employ reconfigurable processing architectures and elements, such as Coarse-Grained Reconfigurable (CGR) Processors (CGRPs) to execute dataflow and/or data parallel computing applications. Accordingly, the disclosure further relates to program compilers of a CGR computing system (CGRS) and compiler allocation of CGRS hardware resources to improve operational efficiency of dataflow and data parallel application programs.

SUMMARY

(3) A method comprises a computer-implemented analysis assistant initiating a presentation of inefficiency results associated with a mapping of operators of a dataflow program to execute on hardware of a computing system to execute the dataflow program. An efficiency analyzer determines the inefficiency results. The assistant initiates the presentation session in response to an interface of a computing system that includes the assistant. In the method, the assistant receives an inefficiency included among the inefficiency results and composes formatted inefficiency results. The formatted inefficiency results comprise a presentation format of the inefficiency to assist a developer of the dataflow program to interpret the inefficiency. The analysis assistant outputs the formatted inefficiency results to an interface of a computing system, and the interface can comprise an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency. In implementations the presentation can comprise an interactive presentation with a developer of the dataflow program.

(4) A computer program product and a computing system can implement the method. The computing system can include a processor to execute the analysis assistant and a processor to execute the efficiency analyzer. The computing system can include the interface to initiate the presentation and/or the interface to output the formatted inefficiency results.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The drawings included in the present disclosure are incorporated into, and form part of, the specification. They illustrate implementations of the present disclosure (hereinafter, “the disclosure”) and, along with the description, serve to explain the principles of the disclosure. The drawings are intended to be only illustrative of certain implementations and are not intended to limit the disclosure.

(2) FIG. 1 illustrates an example coarse-grain reconfigurable (CGR) system (CGRS), according to

aspects of the disclosure.

(3) FIG. 2 illustrates an example sub-graph, according to aspects of the disclosure.

(4) FIG. 3 illustrates an example compiler stack, according to aspects of the disclosure.

(5) FIG. 4A illustrates an example mapping decision space, according to aspects of the disclosure.

(6) FIG. 4B illustrates an example structure of an application analyzer and compiler, according to aspects of the disclosure.

(7) FIG. 5 illustrates an example graph comprising pipelines, according to aspects of the disclosure.

(8) FIG. 6 illustrates an example compiler, according to aspects of the disclosure.

(9) FIG. 7 illustrates an example method to try alternative compilations, according to aspects of the disclosure.

(10) FIG. 8 illustrates an example system for performing trial compilations of a dataflow application, and/or operations thereof, according to aspects of the disclosure.

(11) FIG. 9 illustrates an example method to determine efficiency of mapping decisions, according to aspects of the disclosure.

(12) FIG. 10 illustrates an example method to present mapping suggestions, according to aspects of the disclosure.

(13) FIG. 11 illustrates an example system for compiling a dataflow application and analyzing compilation results, according to aspects of the disclosure.

DETAILED DESCRIPTION

(14) Aspects of the present disclosure (hereinafter, “the disclosure”) relate to methods of compiling neural network applications for execution on computing systems utilizing reconfigurable dataflow processing elements, in particular utilizing coarse-grain reconfigurable processors (CGRPs). More particular aspects relate to determining mappings of neural network operators and data flow to CGRP processing and/or memory elements, and/or configurations of CGRP processing and/or memory elements. Implementations of the disclosure (hereinafter, “implementations”) can analyze a computation graph of a machine learning application or model to determine alternative mappings.

(15) Processing elements that implement aspects of the disclosure can include processors of data parallel (DP) and/or dataflow computing systems, such as Central Processing Unit (CPUs), Graphics Processing Units (GPUs), Field Programmable Gate Arrays (FPGAs), and Digital Signal Processors (DSPs). Certain aspects of the disclosure relate to executing neural networks on computing systems utilizing reconfigurable processors, such as CGRPs, GPUs, FPGAs, reconfigurable Application Specific Integrated Circuits (ASICs), and/or Application Specific Instruction-set Processors (ASIP).

(16) Implementations that are not mutually exclusive are taught to be combinable. One or more features of an implementation can be combined with other implementations. The disclosure in some instances repeats references to these options. However, omission from some implementations of recitations that repeat these options should not be taken as limiting the combinations taught in the preceding sections—these recitations are hereby incorporated forward by reference into each of the following implementations.

(17) Particular expressions of the disclosure will be understood to have the following operative meanings: The phrases “at least one”; “one or more”; and “and/or” are to be understood as open-ended expressions that operate both conjunctively and disjunctively. For example, each of the expressions “at least one of A, B, and C”, “at least one of A, B, or C”, “one or more of A, B, and C”, “one or more of A, B, or C”, and “one or more of A, B, and/or C” means A alone, B alone, C alone, A and B together, A and C together, B and C together, or A, B, and C together. The term “a” or “an” entity refers to one or more of that entity. As such, the terms “a”/“an”, “one or more”, and “at least one” can be used interchangeably herein. The terms “comprising”, “including”, and “having” can be used interchangeably herein.

(18) Unless otherwise specified, the use of ordinal adjectives first, second, third, etc., to describe an object, merely refers to different instances or classes of the object and does not imply any ranking

or sequence.

(19) As used herein, “incorporated subject matter” refers, collectively, to subject matter disclosed, and/or otherwise encompassed, among the disclosures incorporated herein by reference. For purposes of illustrating the disclosure, but not intended to limit implementations, various terms of the disclosure are drawn from the incorporated subject matter. As used herein, unless expressly stated otherwise, such terms as can be found in the incorporated subject matter have the same meanings, herein, as their meanings in their respective incorporated disclosures.

(20) Aspects of the disclosure can be appreciated through a discussion of example implementations and/or applications of methods and/or systems. However, such examples are for purposes of illustrating the disclosure. It should be understood that the intention is not to limit the disclosure to the example implementations described herein, but to encompass all modifications, equivalents, and alternatives falling within the spirit and scope of the disclosure. Thus, the disclosure is not intended to be limited to the implementations shown but is to be accorded the widest scope consistent with the principles and features disclosed herein. Various modifications to the disclosed examples will be readily appreciated by those of ordinary skill in the art, and the general principles defined herein can be applied to other implementations of the disclosure without departing from the spirit and scope of the disclosure.

(21) The disclosure uses terms and acronyms related to the field of the technology, defined, at least in part, herein as: AI—artificial intelligence. AIR—arithmetic or algebraic intermediate representation. ALN—array-level network. Application Model—In machine learning applications, “application model” commonly refers to a mathematical representation of a machine learning application. An application model can comprise an application graph and/or textual (e.g., high level, intermediate level, and/or low level programming language) representation. An application model can represent a set of mathematical operators (compute functions of an application) and a flow of data between the operators, and can represent the operators and dataflow graphically and/or textually. As used herein, “application model” or, simply, “model” refers interchangeably to an application itself (e.g., high level programming statements of an application) and a graphical and/or textual representation of the application's compute functions and/or dataflow. Buffer—an intermediate storage of data. CGR—coarse-grained reconfigurable. A property of, for example, a system, a processor, an architecture (see CGRA), an array, or a unit in an array. This property distinguishes the system, etc., from field-programmable gate arrays (FPGAs), which can implement digital circuits at the gate level and are therefore fine-grained configurable. CGRA—coarse-grained reconfigurable architecture. A data processor architecture that includes one or more arrays (CGR arrays) of CGR units. CGR unit—a circuit that can be configured and reconfigured to locally store data (e.g., a memory unit or a partition memory unit, such as described in Prabhakar), or to execute a programmable function (e.g., a processor or other compute unit, or a partition compute unit such as described in Prabhakar). A CGR unit includes hardwired functionality that performs a limited number of functions used in computation graphs and dataflow graphs. Some implementations include switches to route data among CGR units. CGR Array—an array of CGR units, coupled with each other through an array-level network (ALN), and coupled with external elements via a top-level network (TLN). In implementations a CGR array can physically implement the nodes and edges of a computation and/or dataflow graph. CGRP—Coarse-grain reconfigurable processor. As used herein, CGRP refers to a processor, or processing element, based on a CGRA—such as an integrated circuit, chip, or module based on, or incorporating, a CGRA—and/or incorporates a CGR unit, CGR array, or elements of a CGR unit and/or a CGR array. CGR Components—As used herein, “CGR components” refers, collectively, to hardware resources or elements of CGR units, CGR arrays, and CGRP; memories of CGR units/arrays/processors; and, networks and/or I/O interconnections and interface hardware interconnecting CGR units/arrays/processors and/or memories, such as Ethernet networks/interfaces, I/O buses/interfaces, such as PCI-Express buses, InfiniBand buses/interfaces, and/or memory or data buses/interfaces, such as buses of a processor

and/or memory fabric, and related interface hardware). CGR hardware—As used herein, the terms “CGR hardware” and “CGR hardware resources” refer to any individual hardware element, or combination of hardware elements, of CGR components of a CGRS. CGRS—a computing system comprising CGR units and/or CGRPs. As used herein, CGRS refers to a computing system that is based on, and/or can utilize, reconfigurable computing resources, such as CGR arrays, CGR units, and/or CGRPs, to perform operations of data parallel and/or dataflow applications. U.S. Nonprovisional patent application Ser. No. 16/239,252, “VIRTUALIZATION OF A RECONFIGURABLE DATA PROCESSOR”, to Grohoski, et al, (hereinafter, “Grohoski”), and U.S. Nonprovisional patent application Ser. No. 16/922,975, “RUNTIME VIRTUALIZATION OF RECONFIGURABLE DATA FLOW RESOURCES”, to Kumar, et al, (hereinafter, “Kumar”), both incorporated herein by reference, illustrate example implementations of CGR arrays, CGR units, CGRPs, and CGR systems. Chip—As used herein, the term “chip” refers to an IC (or, combination of ICs) that can embody elements of a CGRA. A chip can typically be packaged in a chip module (e.g., a single chip module, “SCM” or, alternatively, a multi-chip module, “MCM”). Compiler—a translator that processes statements written in a programming language to machine language instructions for a computer processor. A compiler can include multiple stages to operate in multiple operations. Each stage can create or update an intermediate representation (IR) of the translated statements. Compiler stages are illustrated with reference to FIG. 3. Computation graph/Graph—As used herein, computation graph refers to a type of directed graph comprising nodes and edges connecting the nodes, to represent a dataflow application. In a neural network application nodes can represent mathematical operations/expressions and edges that indicate dependencies between the operations/expressions. For example, in machine learning (ML) algorithms, input layer nodes can assign variables, output layer nodes can represent algorithm outcomes, and hidden layer nodes can perform operations on the variables. Edges can represent data (e.g., scalars, vectors, tensors) flowing between operations. In addition to dependencies, the computation graph reveals which operations and/or expressions can be executed concurrently. Dataflow Application—As used herein, the term “dataflow” application refers interchangeably to data parallel and dataflow applications. such as ML, AI, and other massively parallel computing applications. Dataflow Graph—a computation graph, or portion of a computation graph, corresponding to operators (application compute functions), data, and flow of data among operators, of a dataflow application that includes one or more loops of operator nodes that can be nested, and wherein nodes can send messages to nodes in earlier (predecessor) layers to control the dataflow between the layers. IC—integrated circuit—a monolithically integrated circuit, i.e., a single semiconductor die which can be delivered as a bare die or as a packaged circuit. For the purposes of this document, the term integrated circuit also includes packaged circuits that include multiple semiconductor dies, stacked dies, or multiple-die substrates. Such constructions are now common in the industry, produced by the same supply chains, and for the average user often indistinguishable from monolithic circuits. Intermediate Representation (IR)—an Intermediate Representation is a representation of an application in an intermediate language. An IR can incorporate partial compilation results, such as sections (groupings) of a graph or model, pipelines that can be formed within a graph or model, mappings of application functions or graph nodes/edges to hardware resources of a CGRS. Logical CGR—A logical CGR array or logical CGR unit comprises a representation of a CGR array or a CGR unit that is physically realizable, but that may not, at a particular time in executing a dataflow application, have been assigned to a physical CGR array or to a physical CGR unit on an IC. ML—machine learning. PEF—processor-executable format—a file format suitable for configuring a configurable data processor. Pipeline—a staggered flow of computational operations through a chain of pipeline stages in which the operations can be executed in parallel. In an application graph, a pipeline can comprise a set of operator nodes that can pipeline operations of the graph. Pipeline Stages—a pipeline can be divided into stages that are coupled with one another as predecessor/successor stage to form a pipe topology. PNR—place and route—the assignment of logical CGR units and

associated processing/operations to physical CGR units in an array, and the configuration of communication paths between the physical CGR units. RAIL—reconfigurable unit abstract intermediate language. RP—reconfigurable processor. An RP can comprise, for example, field programmable gate arrays (FPGAs), graphic processing units (GPUs), and/or CGRPs. TLIR—template library intermediate representation (IR). TLN—top-level network.

(22) Turning now to more particular aspects of the disclosure, high-level programs for machine learning (ML) and artificial intelligence (AI) can require massively parallel computations, where many parallel and interdependent computation threads (pipelines) exchange data. Such programs are ill-suited for execution on traditional, Von Neumann architecture computers. Rather, these applications can require architectures optimized for parallel and pipeline processing, such as CGRAs or graphic processing units (GPUs).

(23) The ascent of dataflow applications such as ML and AI, and massively parallel architectures (such as CGRAs) places new and complex requirements to execute the applications, or computations of the applications, on CGR hardware. Such requirements can include how computations of an application are pipelined, which computations are assigned to which compute units, how data is routed between various compute units and memories, and how synchronization among processors, memories, and data transfer hardware is controlled, particularly when a dataflow applications includes one or more nested loops, whose execution time can varies depending on the data being processed. The architecture, configurability and dataflow capabilities of CGR systems, and CGR components of CGR systems, enable increased compute power that supports both parallel and pipelined computation.

(24) In implementations CGR components of a CGRS, for example, can be programmed to simultaneously execute multiple independent and interdependent operations. To enable simultaneous execution within a pipeline stage, and across pipeline stages, dataflow applications need to be distilled from a high-level program and translated to low level instructions to execute the program on hardware resources of reconfigurable dataflow systems, such as a CGRS. The low level instructions can comprise a configuration file describing a configuration of CGR components, as well as processor (e.g., CGRP) instructions and/or instructions for transferring application data among CGR components.

(25) A high-level program is source code written in programming languages like Spatial, Python, C++, and C, and can use computation libraries for scientific computing, ML, AI, and the like. The high-level program and referenced libraries can implement computing structures and algorithms of machine learning models like AlexNet, VGG Net, GoogleNet, ResNet, ResNeXt, RCNN, YOLO, SqueezeNet, SegNet, GAN, BERT, ELMo, USE, Transformer, and Transformer-XL.

(26) In computing applications, a compiler translates high-level programs to instruction executable by processors of a computing system. In a CGRS, a CGRS compiler can translate high-level programs to processor instructions, but also to executable instruction files and/or “bit files” describing configurations of CGR components to execute a dataflow application, or pipeline stages of a dataflow application. CGRS compilers require mapping application operations and data flow to CGR hardware components in both space (CGR hardware parallelism) and time (for synchronization of interdependent computations). This requirement implies that a CGRS compiler must determine which operations of a dataflow application are assigned to which of the CGR components, and how both data and, related to the support of computation and control information flow among CGR components, and to/from external hosts and storage. This process, known as “place and route”, is one of many new challenges posed to CGRS compilers.

(27) FIG. 1 illustrates an example reconfigurable dataflow computing system **100** including a CGR processor **110**, a host **180**, and a memory **190**. CGR processor **110** has a coarse-grained reconfigurable architecture (CGRA) and includes an array of CGR units **120** such as a CGR array. CGR processor **110** further includes an IO interface **138**, and a memory interface **139**. Array of CGR units **120** is coupled with IO interface **138** and memory interface **139** via data bus **130** which

can be part of a top-level network (TLN). Host **180** communicates with IO interface **138** via system data bus **185**, and memory interface **139** communicates with memory **190** via memory bus **195**. Array of CGR units **120** can further include compute units and memory units that connected with an array-level network (ALN) to provide the circuitry for execution of a computation graph or a dataflow graph that can have been derived from a high-level program with user algorithms and functions. The high-level program can include a set of procedures, such as learning or inferencing in an AI or ML system. More specifically, the high-level program can include applications, graphs, application graphs, user applications, computation graphs, control flow graphs, dataflow graphs, models, deep learning applications, deep learning neural networks, programs, program images, jobs, tasks and/or any other procedures and functions that can need serial and/or parallel processing. In some implementations, execution of the graph(s) can involve using multiple units of CGR processor **110**. In some implementations, CGR processor **110** can include one or more ICs. In other implementations, a single IC can span multiple CGR processors. In further implementations, CGR processor **110** can include one or more units of array of CGR units **120**.

(28) Host **180** can be, or can include, a computer such as will be further described with reference to FIG. **11**. Host **180** can execute runtime processes, as further referenced herein, and can also be used to run computer programs, such as a CGRS compiler. In some implementations, the compiler can run on a computer that is similar to the computer described with reference to FIG. **11**, but separate from host **180**.

(29) CGR processor **110** can accomplish computational tasks by executing a configuration file (for example, a PEF file). For the purposes of this description, a configuration file corresponds to a dataflow graph, or a translation of a dataflow graph, and can further include initialization data. A compiler compiles the high-level program to provide the configuration file. In some implementations described herein, a CGR array is configured by programming one or more configuration stores with all or parts of the configuration file. A single configuration store can be at the level of the CGR processor or the CGR array, or a CGR unit can include an individual configuration store. The configuration file can include configuration data for the CGR array and CGR units in the CGR array, and link the computation graph to the CGR array. Execution of the configuration file by CGR processor **110** causes the CGR array (s) to implement the user algorithms and functions in the dataflow graph.

(30) CGR processor **110** can be implemented on a single integrated circuit die or on a multichip module (MCM). An IC can be packaged in a single chip module or a multichip module. An MCM is an electronic package that can comprise multiple IC dies and other devices, assembled into a single module as if it were a single device. The various dies of an MCM can be mounted on a substrate, and the bare dies of the substrate are electrically coupled to the surface or to each other using for some examples, wire bonding, tape bonding or flip-chip bonding.

(31) Many dataflow applications, such as in ML and other types of AI applications, comprise neural networks (NNs). Examples of neural networks include fully connected neural networks (FCNNs), recurrent neural networks (RNNs), graph neural networks (GNNs), convolutional neural networks (CVNNs), graph convolutional networks (GCNs), long short-term memory (LSTM) networks, autoencoders, deep belief networks, and generative adversarial networks (GANs).

(32) In data parallel and dataflow applications, such as NNs, compute functions of the application are often referred to as “operators”. The compute functions perform computations, such as matrix computations using tensor data of the application, to execute the higher level processes of the application (e.g., object recognition in an image, natural language phrase interpretations or prediction, etc.). A neural network processes data according to a flow of computational input (operand) and computational output (results) data through layers of operators (neurons) of the NN.

(33) Operators of an input layer can receive stimuli (e.g., input data), and the input and other (e.g., “hidden”) layers compute particular functions (e.g., an activation or loss function), and operators of an output layer output computational results. A particular layer of an NN comprises operators that

perform the particular function computations of that layer. Example layers, and associated operators, of NNs include rectified linear unit (ReLU) layers, fully connected layers, recurrent layers, graphical network layers, long short-term memory layers, convolutional layers, kernel layers, dropout layers, and pooling layers.

(34) A machine learning application requires “training” within a problem space the application is designed to recognize (e.g., subjects of images, audio, or video) or predict outcomes (e.g., natural language phrase completion, future values, etc.). Training a neural network can comprise determining and/or optimizing parameters associated with computations (e.g., activation functions) of the NN computed by operators within layers of the NN. Weights and biases, for example, can be parameters of a weights-bias activation function of a neural network. In training such an NN, a training (data parallel/dataflow) application can compute gradients of weights and biases, such as by using a loss-function, and can optimize the weights and biases based on an optimization algorithm such as gradient descent. Executing an ML application can utilize the optimized parameters to execute functions of the application.

(35) Problem spaces of a machine learning application, and/or input of dataflow applications in general, can comprise enormous amounts of data, and can often comprise tensor data. Thus, functions of these applications (e.g., operators of neural networks) commonly involve linear algebra computations over tensor data, such as matrix multiplication, transposition, and addition. Algorithms commonly employed in dataflow applications include algorithms such as linear regression and gradient descent over tensors and/or matrices of tensors. Matrices of tensors data can comprise matrices of varying dimensions and a variety of computing systems, including dataflow computing systems, can perform matrix computations, such as GeMM, matrix summation, matrix transposition, gradient computations, and/or backpropagation of matrix computations, to process tensors in dataflow applications such as machine learning in neural networks.

(36) As used herein, brackets and a capital letter, such as [M], is used to refer to a matrix as a whole, while lowercase letters, such as m, are used to refer to an element, or set of elements, of a matrix [M]. For example, an expression such as ($w \times a$) refers, herein, to a multiplication of a set of elements of matrices [W] and [A], such as elements of a row of matrix [W] multiplied by elements of a corresponding column of matrix [A]. The term “element”, in reference herein to a matrix, refers to the contents (e.g., a scalar value) of a row and column cell of the matrix.

(37) A common computation for processing tensors in dataflow applications is a sum of products (dot product) of two matrices. The products comprise products of elements of a row of one multiplicand matrix (a “left side” matrix multiplied by corresponding elements of a column of a second multiplicand (a “right side” matrix), where the row dimension of the left side matrix and the column dimension of the right side are the same (shared dimension.) As used herein, the term “dot product” refers to a sum of two or more products of a row of a left side matrix multiplicand by a column of a right side matrix. An expression such as ($\sum w a$) refers to a sum-product of elements w and a (e.g., a sum of products $w \times a$ for elements of a row of a matrix [W] multiplied by elements of a column of a matrix [A]). As an example, a dot product of elements w.sub.11 of matrix [W] multiplied by a.sub.11 of matrix [A], and w.sub.11 multiplied by a.sub.21 of matrix [A], is [$w_{sub.11} \times a_{sub.11} + w_{sub.11} \times a_{sub.21}$].

(38) A “matrix summation” computation, as used herein, refers to a matrix computation in which a dot product of two multiplicand matrices is added to a matrix addend. A matrix addend can comprise a constant or can comprise a matrix (which can itself be multiplied by a matrix multiplied by a constant) sharing a row dimension of the dot product of two multiplicand matrices. A “weight-bias function”, $y = \sum w a + b$, is one example of such a computation, in which a weights matrix [W] is multiplied by an activation matrix [A] and the dot products, $\sum w a$, for each row/column set of products, is added to elements of a bias matrix [B].

(39) In implementations, a CGRP, and/or other CGR components of a CGRS, can perform

computations (e.g., operators) of applications in a distributed fashion and/or can execute computations as dataflow pipelines that can efficiently exploit CGRS and application parallelism, and CGR component data locality. Dataflow pipelines of CGRS compute units (e.g., CGRPs and/or CGR arrays) can contain several computational stages, in which each stage can read data from one or more input buffers (e.g., buffers in CGR component memories), can perform computations on the data while using one or more internal buffers to store and retrieve intermediate results, can produce outputs, and can write the outputs to one or more output buffers.

(40) Data parallel and dataflow computing applications can comprise tensor computations, usually involving enormous amounts of data, such as very large and/or numerous matrices of tensor data. For example, machine learning (ML) and other tensor-based applications can comprise a convolutional neural network (NN). While not intended to limit implementations, a convolutional neural network can serve to illustrate aspects of the disclosure. However, it will be appreciated by one of ordinary skill in the art that aspects of the disclosure can apply broadly to a variety of computing applications involving tensor data, and/or executed by data parallel and/or dataflow applications and computing systems.

(41) An NN can comprise layers organized as a pipeline of computations using matrices of tensor data. A layer of the NN can comprise operators performing computations on matrices of tensor data. A particular operator of an NN (or, tensor-based application in general) can perform a matrix computation, such as Generalized Matrix Multiplication (“GeMM”), matrix convolution, and Rectified Linear Units (“ReLU”) corresponding to particular algorithms and/or functions of the application, such as an activation function, gradient descent function, and/or a loss function. A particular layer of an NN can comprise multiple processing elements, such as CGRPs, executing in parallel to perform operator computations of the application using subsets of tensor data. The processing elements of one layer of an NN can output results of their computations to a successor “forward” and/or “backward” layer of the NN.

(42) Various types and/or combinations of computing systems can execute tensor-based applications, and/or operators of tensor-based applications, such as NNs. Data parallel (DP) and dataflow computing systems, particularly systems utilizing CGRPs, can be particularly efficient at executing tensor-based applications. CGRPs can individually, or in combination, execute functions and/or computations of application operators, in parallel and in pipelines, to efficiently execute an application and improve performance of application execution. As used herein, the term “reconfigurable dataflow system (DS)” refers, interchangeably, to data parallel and dataflow computing systems utilizing reconfigurable processors such as CGRPs. A CGRS can, for example, efficiently execute tensor-based applications such as convolutional neural networks, and can serve to illustrate aspects of the disclosure without limiting implementations.

(43) A tensor-based application can include “operators” that perform computations such as linear regression, non-linear regression, Gaussian regression, Support Vector Machine (SVM) regression, Generalized Linear Models, regression trees, shallow and deep neural network models, logistic regression, decision tree, and, “K” nearest neighbor, using matrices of tensor data. One expression, or representation, of an application is a computation graph (hereinafter, for brevity, simply “graph”), which can be textual, graphical, or a combination of textual and graphical descriptions of operators, operands, and results of computations of the application. A graph can represent the operators (as compute nodes of the graph) of an application, and their arrangement and/or dependencies (e.g., flow of computational inputs and outputs) among the operators (as edges of the graph).

(44) Data nodes of a graph can represent particular application data elements, such as input data for training an ML model. A graph can be a directed acyclic graph (DAG), or can comprise loops, and even nested loops, of operators. As used herein, except where otherwise qualified as “data node”, the term “node” is used herein interchangeably to refer to an operator of an application and a node representation of that operator in a graph.

(45) Forward nodes of a graph can receive outputs of backward nodes (e.g., gradients), and backward nodes can receive updated outputs of forward nodes (e.g., outputs computed using outputs of backward nodes), creating feedback loops within the graph. As nodes within a feedback loop recompute outputs based on the feedback, such nodes are referred to herein as “recompute nodes”.

(46) A pipeline of an application can comprise a set of forward operators and, optionally, set of backward operators (e.g., backpropagation operators). Each operator within a pipeline can process data output from a predecessor operator, generally in parallel with the predecessor operators as the predecessor operator outputs results of computations over a portion input data.

(47) FIG. 2 illustrates an example of a computation graph corresponding to an application. As shown in FIG. 2, forward and backward operators of an application can be grouped, such as for mapping the operators to CGR components for execution, as respective forward and backward sections of a graph. The sections can each represent nodes of the graph that do not have data dependencies among each other (that is, do not need to await complete computational results of another compute node), such that a CGRS can execute computations of the nodes in a pipeline topology among CGR components. Sections can particularly comprise operators that are “adjacent” and, based on not having data dependencies among each other, can form a pipeline. A “producer” operator and a “consumer” operator are adjacent operators in a graph if the producer can input results of the producer operator computations to an operand of the consumer operator. For example, a GeMM operator is adjacent to an ADD operator if a results matrix (or elements of a results matrix) of the GeMM operator can be direct input to the ADD operator.

(48) In FIG. 2, forward sections **210** is shown comprising Pipe **214A** and Pipe **214B**, and backward sections **220** is shown comprising Pipe **224A** and Pipe **224B**. Pipe **214A** is shown comprising node CONV **212A**, and Pipe **224B** is shown comprising nodes RELU **212B**, CONV **212C**, RELU **212D**, and MAXPOOL **212E** (hereinafter, collectively “nodes **212**”). Names of nodes, such as a “RELU”, can indicate a type of computation of the application performed by a node.

(49) Edges of a graph can represent data flow between and into or out of the nodes. Thus, computational results of node CONV **212A** can flow as inputs to node RELU **212B**, computational results of node RELU **212B** can flow as inputs to node CONV **212C**, and so forth. Data nodes in a graph can represent data processed by compute nodes and flow of data into or out of the nodes (as also shown in FIG. 2 by directed arrows). In forward sections **210**, FIG. 2 depicts data nodes OP DATA **202** and WEIGHT **204** as data input to CONV **212A**, and WEIGHT **206** as data input to CONV **212C**.

(50) In FIG. 2, backward sections **220** is shown comprising Pipe **224A** and Pipe **224B**, Pipe **224A** is shown comprising nodes CONV2D BWD **222A** and RELU BWD **222B**, and Pipe **224B** is shown comprising nodes CONV2D BWD **222C**, RELU BWD **222D**, and MAXPOOL **222E**. In backward sections **220**, FIG. 2 depicts WEIGHT **206** as data input also to CONV2D BWD **222C**. Backward nodes of a graph can represent nodes that receive outputs of forward nodes and compute a feedback function over those outputs. For example, a common backward computation is to compute gradients of weights and biases, and/or loss functions based on gradients of weights and biases, in a weights-bias activation function of a forward node. Backward nodes, can compute, for example, a gradient in an application that include gradient descent to optimize computations of forward nodes in a feedback loop. As shown, an output of backward sections **220** is data node output gradient **208**, output node CONV2D BWD **222A**.

(51) In implementations, a “CGRS compiler” can compile a high-level language representing of a data parallel and/or dataflow application to configurations and/or execution instructions to execute the application. For brevity, hereinafter “application” is understood to refer to a data parallel or dataflow programming application for execution by a data parallel and/or dataflow computing system, such as a CGRS.

(52) A CGRS compiler can, for example, transform an application into, and/or can utilize, a graph

such as example graph **200** in FIG. 2. Based on a graph of an application, a CGRS compiler can generate a search space, and can use the graph and/or search space to determine application operational parallelism and pipelining, and/or to map application dataflow (e.g., nodes and edges of a computation graph) to CGRS and/or CGR hardware resources and dataflow through the resources. A compiler can further transform resource mapping decisions into assembler input for generation of hardware instructions and/or hardware configuration files, such as a Processor Executable Format (PEF) file.

(53) FIG. 3 is a block diagram of example compiler stack **300** comprising multiple compilation stages to compile a dataflow application for execution by a CGRS. As depicted in FIG. 3, compiler stack **300** includes several stages to translate a high-level program, with (user) dataflow application algorithms and functions (e.g., ML algorithms and/or tensor computation functions), to configuration and/or instruction data for a CGRS to execute the application.

(54) Compiler stack **300** can take its input from application platform **310**, and/or any other source of high-level program statements of an application, which provides a user interface, such as an API and/or command line interface (CLI), for application developers to compile an application. A “user”, as used herein, can be any human or computing system that develops an application (e.g., programs the high-level programs of an application), and/or that can input an application into a CGRS compiler for translation to CGRS configurations and/or CGRS execution instructions.

(55) Compiler stack **300** can further receive hardware description **315**, which can comprise a textual and/or graphical description of CGRS and/or CGR hardware components of a CGRS. Compiler stack **300** can utilize hardware description **315** to translate the high-level programming statements of an application to configurations CGR components and/or execution instructions (e.g., instructions to a runtime processor to control execution, and/or processor instructions to execute functions, of an application) to execute the application.

(56) Application platform **310** can comprise a computing system for developing an application and/or inputting an application for compilation by a CGRS compiler. For example, application platform **310** can comprise a computing system capable of hosting a user, such as host processor in the CGRS examples of Kumar. Application platform **310** can include libraries such as PyTorch, TensorFlow, ONNX, Caffe, and Keras to provide user-selected and configured algorithms.

(57) Application platform **310** can output a high-level program of an application to compiler **320**, which in turn can output a configuration file to runtime processes **330**. Runtime processes **330** can comprise programs to configure CGR components, and/or manage execution of an application on CGR components, of a CGRS. The programs can execute on a runtime processor (e.g., one or more CPUs) of a CGRS.

(58) Compiler **320** can include dataflow graph compiler **321**, algebraic graph compiler **322**, template graph compiler **323**, template library **324**, and placer and router PNR **325**. In implementations, template library **324** can include a reconfigurable unit abstract intermediate language (RAIL), and/or assembly language interfaces (APIs) for power users.

(59) Dataflow graph compiler **321** can analyze high-level programs, implementing user algorithms and application functions received from application platform **310**, and can convert the high-level programs to one or more dataflow graphs. The high-level programs can be suitable for parallel and/or pipeline processing and nodes of the dataflow graphs can be intrinsically parallel unless an edge in the graph indicates a dependency. Dataflow graph compiler **321** can provide code optimization operations, such as false data dependency elimination, dead-code elimination, and numeric constant folding. The dataflow graphs can encode data and execution control dependencies of the high-level programs.

(60) Dataflow graph compiler **321** can support programming a CGR components (e.g., CGRPs) using higher or lower-level programming languages. For example dataflow graph compiler **321** can support translation or conversion from an application platform **310** to C++ and/or an assembly language. In implementations, dataflow graph compiler **321** can allow programmers to provide

code (e.g., machine language code) that runs directly on CGRPs and/or other CGR components. Dataflow graph compiler **321** can include one or more programming libraries, and the libraries can include predefined functions, such as linear algebra operations, element-wise tensor operations, non-linear functions, and reduction functions for creating, executing, and profiling dataflow graphs on the CGRPs. Via the application platform **310**, dataflow graph compiler **321** can provide an API to enhance programming functionality available to application developers.

(61) Algebraic graph compiler **322** can include an Application analyzer and Compiler (MAC) level that can make high-level mapping decisions for sub-graphs (also referred to as “sections” or “section cuts”) of a dataflow graph based on CGR hardware constraints. Algebraic graph compiler **322** can support various application frontends, such as Samba, JAX, and TensorFlow/HLO. Algebraic graph compiler **322** can also transform the graphs, for example via autodiff and GradNorm, to perform stitching between sub-graphs, interface with template generators for performance and latency estimation, convert dataflow graph operations to algebraic intermediate representation (AIR) operations, perform tiling, sharding (database partitioning) and other application preparation operations, and can model or estimate execution parallelism that can be achieved within the dataflow graphs.

(62) Algebraic graph compiler **322** can include an arithmetic or algebraic intermediate representation (AIR) level that can translates high-level dataflow graph and mapping decisions provided by a MAC level into AIR graphs. An AIR level can include validating and/or correcting (“legalizing”) a dataflow graph and/or mapping decisions of a MAC; expanding data parallel, tiling, pipeline, and/or region instructions provided by a MAC; inserting stage buffers and skip buffers, eliminating redundant operations, buffers, and sections; and, optimizing resource use, execution latencies, and computational throughput.

(63) Template graph compiler **323** can translate AIR graphs to a template library intermediate representation (TLIR). A TLIR can comprise a graph that can optimize configurations and/or execution instructions based on target (CGRS and/or CGR) hardware architecture and/or to unplaced units suitable for place, allocate, and route level PNR **325**. Template graph compiler **323** can add further information node names, node inputs, node input names, and dataflow descriptions) as inputs to PNR **325**, and can make the graph physically realizable through each layer of the graph. Template graph compiler **323** can, for example, translate AIR graphs to specific application operation templates, such as templates for general matrix multiplication (GeMM), matrix transposition, and/or matrix convolution operations. In implementations a CGRS compiler like compiler **320** a can convert part or all intermediate representation operations to templates, stitch templates into data and control flow of the application, insert necessary buffers and layout transforms, generate test data, and optimize for CGR hardware utilization, execution latency, and compute and/or data transfer throughput.

(64) Implementations can use templates for common operations. Templates can be implemented using assembly language, RAIL, or similar language and/or representation constructs. RAIL can compare to a low-level language, in that memory units and compute units can be separately programmed in RAIL constructs, but RAIL can provide a higher level of abstraction and compiler intelligence that, for example, an assembly language, via a concise performance-oriented and domain-specific language for CGR component (e.g., CGR array) templates. RAIL can enable template writers and external power users to control interactions between logical compute units and memory units of CGR components using high-level expressions, without the need to manually program actions such as capacity splitting, register allocation, etc. RAIL logical compute and memory units can also enable stage/register allocation, context splitting, transpose slotting, resource virtualization and mapping to multiple physical compute units and memory units (e.g., PCUs and PMUs of tiles, such as in the examples of Grohoski and Kumar).

(65) Template library **324** can include an assembler that provides an architecture-independent, low-level programming interface as well as optimization and code generation for CGR hardware. An

assembler can include memory address expression compilation, CGR hardware intra-unit resource allocation and management, rendering a template graph physically realizable based on CGR hardware-specific rules, low-level CGR hardware architecture-specific transformations and optimizations, and CGR hardware architecture-specific code generation.

(66) PNR **325** can translate RAIL and/or assembly language outputs of template library **324**, and/or TLIR outputs from template graph compiler **323**, and can map logical (e.g., unplaced physically realizable) CGR units, to physical CGR hardware implementation levels, such as an SCM, MCM, and/or chip level of CGR components, can determine physical data channels to allow for communication among the CGR units and between the CGR components (e.g., components coupled via a TLN, allocate memory, I/O, and/or switch ports of CGR components, provide CGR component configuration data and initialization data, and can produce configuration files, e.g., processor-executable format (PEF) files. PNR **325** can provide bandwidth calculations, allocate network interfaces, provide configuration data for CGR components to perform memory address translation, and control switch and data routing among CGR components. PNR **325** can perform such functions in multiple operations and can include multiple modules (not shown in FIG. 3) to perform the multiple operations, e.g., a placer, a router, a port allocator, and a PEF file generator). PNR **325** can receive input data, for example, from any of the higher-level modules (dataflow graph compiler **321**, algebraic graph compiler **322**, template graph compiler **323**, and/or template library **324**). In implementations, a higher-level module, such as template graph compiler **323**, can prepare information for PNR **325** and can omit other levels directly providing input data to PNR **325**.

(67) Implementations of compiler **320** compile applications in an iterative process, such as feeding information from PNR **325** back to a higher-level module, which can, in turn, execute a new compilation operation using physically realized results, rather than estimates of, or logical placeholders for, physically realizable circuits. For example, PNR **325** can feed information regarding the physically realized circuits back to algebraic graph compiler **322**.

(68) Memory allocations can represent logical memory spaces in on-chip (a chip implementing a CGR component) and/or off-chip (a chip separate from a CGR component), CGR component memories, for data flowing through the dataflow graph; a configuration file, such as a PEF, can specify particular memory allocations. Memory allocations can define a type and number of CGR hardware memories and/or circuits (functional units, storage, or connectivity components). Main memories (e.g., DRAM) can be, for example, off-chip memories, and scratchpad memories (e.g., SRAM) can be on-chip memories, such as memories of a CGR array. Memory allocations can correspond to various access patterns and/or memory layouts, such as access patterns/layout of cache memories, read-only look-up tables (LUTs), serial memories (e.g., FIFOs), and/or register files.

(69) Compiler **320** can bind memory allocations to unplaced memory units and can bind operations of a dataflow graph to unplaced compute units, for execution of a graph, and configuration data, such as in a PEF, can specify such bindings. In implementations, compiler **320** can partition parts of a dataflow graph into memory subgraphs and compute subgraphs, and can specify these subgraphs in configuration file. A memory subgraph can comprise, for example, address calculations leading up to a memory access. A compute subgraph can comprise, for example, compute operations (compute nodes) in a parent graph. A compiler can divide a parent graph into multiple memory subgraphs and a single compute subgraph, for example. A single parent graph can produce one or more memory subgraphs, depending on how many memory accesses exist in the original graph loop body. In cases where the same memory addressing logic is shared across multiple memory accesses, a compiler can duplicate address calculations to create multiple memory subgraphs from the same parent graph.

(70) Compiler **320** can generate configuration files with configuration data (e.g., a bit stream) for the placed positions, and for routed data and control networks. In implementations this can include

the compiler assigning coordinates and communication resources of the physical CGR components by placing and routing unplaced units of CGR components with a goal to maximize compute and/or data transfer bandwidth and minimizing compute and/or data transfer latency.

(71) An application may not itself include backward nodes and, in implementations, a CGRS compiler, such as illustrated by the example of compiler **320**, can determine that an application requires backward nodes, and can generate backward nodes in a computation graph. In determining a mapping of an application to CGR hardware resources, a CGRS compiler can identify recompute nodes and can determine section boundaries among forward nodes, backward nodes, and recompute nodes within a graph.

(72) To exploit the full power of a CGRS—particularly, dynamically reconfigurable CGR components of a CGRS—a CGRS compiler must not only generate low level processor instruction sequences, but must also allocate reconfigurable resources of the underlying CGR hardware that can execute the application most efficiently, and with highest possible computational performance. A CGRS compiler must, further, determine controls to sequence transfer in (e.g., to a memory and/or compute unit), processing (e.g., compute unit and/or operator pipelining), and/or transfer out (e.g., from a memory and/or compute unit) of application data.

(73) In optimizing parallelization and computational latency of among CGRS hardware resources, a CGRS compiler must consider complex factors, such as: the number of available processing units (e.g., processors of CGR components); the number, size, and transfer latency of memory units (e.g., memories of CGR components); computational latency of operators of the application; dependencies among operators; and, sections of an application that can execute in parallel, not only intrinsically, but also given the amount of CGRS hardware resources available to execute the sections.

(74) Such considerations can be referred to as “mapping factors”. In implementations “mapping decision space” can comprise mapping factors. In addition, or alternative, to factors just described, the mapping factors can include parameters and/or attributes of an application and/or CGRS related to mapping factors, such as just described. Mapping factors included in a mapping decision space can include, for example, descriptions and/or attributes of CGR components; configurations and/or arrangements of data nodes, compute nodes, and interconnections of nodes (edges) of a graph and CGR components; and/or, groupings (“section cuts”) of operators of a graph into particular pipelines and sections. Mapping factors of a mapping decision space can include alternative such configurations and section cuts, and can include costs (e.g., hardware utilization, compute and/or data transfer bandwidth or latency) associated with the alternatives. Mapping factors of a mapping decision space can include optimization goals (e.g., optimizing utilization over latency, or vice versa) and/or priorities of execution of particular nodes of a graph.

(75) FIG. 4A illustrates mapping factors and a mapping decision space a CGRS compiler can utilize in mapping operators and data of an application to underlying hardware resources of a CGRS (e.g., CGR components of a CGRS). A MAC component of a CGRS compiler, for example, can generate and/or analyze a computation graph of an application to determine mapping factors of a mapping decision space. For example, a MAC can traverse a graph, such as in the example of FIG. 2, to determine mapping factors of a mapping decision space.

(76) In implementations, a compiler can determine a mapping of an application (e.g., operators and tensors included in a graph of an application) to CGR hardware resources for execution of the application. A compiler, or a MAC of a compiler, can include a hardware mapping component—referred to herein as a “mapper”—and the mapper can analyze a graph to map operators, tensors, and/or tensor dataflow of an application to CGR hardware for execution.

(77) For purpose of illustrating the disclosure, example operations of the disclosure, such as example operations of FIG. 4A, are frequently described as performed by a MAC, and/or components of a MAC, of a CGRS compiler. However, this not intended to limit implementations and one of ordinary skill in the art will appreciate that a compiler need not necessarily comprise a

CGRS compiler, a MAC of a CGRS compiler, and/or particular components (e.g., a mapper”) of a compiler or a MAC to perform methods, and/or operations of methods, of the disclosure. Components of a compiler alternative to these particular components can perform methods and operations of the disclosure within the scope and spirit of the disclosure.

(78) In FIG. 4A, decision space **400** is an example of a mapping decision space that a CGRS compiler can utilize to determine alternatives to map an application to CGR hardware for a CGRS to execute the application efficiently. Decision space **400** can represent a combination (not necessarily exhaustive) of mapping factors **402-412** (collectively, “mapping factors **400**” in FIG. 4A) that a CGRS compiler can include in a mapping decision space such as example decision space **400**.

(79) In FIG. 4A, app **418** can comprise an application, and/or application model, (e.g., represented as a graph and/or textual representation) and MAC **416**, in FIG. 4A, can be a MAC component of a CGRS compiler configured to compile app **418**. MAC **416** can generate decision space **400** to execute app **418** on CGR hardware that can be represented by hardware attributes **414**. In the example of decision space **400**, mapping factors **400** are shown in FIG. 4A including PAR factors **402**, tiling factors **404**, model/data parallelism **406**, stage boundaries **408**, recompute sections **410**, and section/HW boundaries **412**.

(80) PAR factors **402** can comprise, for example, parallelization (“PAR”) factors included in a template (e.g., a template among template library **324** in FIG. 3) that can represent an intrinsic, or application programmer preferred, parallelization of application operators. Tiling factors **404** in decision space **400** can include alternative, and/or optimal, tiling of operator and/or pipeline input data, operand matrices, and/or operator results matrices. Tiling a graph refers to partitioning, or “slicing”, operand/results matrices of tensor data input to, and output from, operators in the graph into smaller matrices (“tiles”). A MAC can tile the matrices based on, and/or to preserve, a particular, shared dimension of the matrices (e.g., a row dimension or a column dimension of the matrices). Application/data parallelism **406** can include boundaries of operator and data parallelism, which can represent, for example, a degree of parallelization of application operators and data. Stage boundaries **408** can include, for example, boundaries of pipeline stages of underlying CGR hardware and/or CGR component hardware.

(81) As illustrated in the examples of FIG. 2, a graph of an application can comprise sections. Operators that cannot be executed in parallel (e.g., operators that cannot be included in a pipeline with another operator) cannot be included in the same section of an application. Similarly, underlying CGR hardware can have limits to the number and/or type of operators that it can perform in parallel, and/or the amount of data it can process (e.g., based on sizes of memory to buffer or store input data and/or computation outputs). Thus, section/HW boundaries **412** can include boundaries, within an application or graph of an application, between forward and backward sections of the application, and/or boundaries of CGR hardware to execute operators within particular sections of a graph. Hardware boundaries among section/HW boundaries **412** can be based on a hardware description, and/or attributes of hardware, of CGR hardware, such as can be included in hardware attributes **414**.

(82) Backward nodes can be feedback paths, in the model, to recompute nodes, and the recompute nodes can be factors of decision space **400**, such as to determine dependencies among sections and operators within sections. Recompute sections **410**, for example, can represent combinations of operators that recompute particular application functions, such as recomputing activation functions using results (e.g., gradient adjusted matrices) of backward section operators.

(83) In implementations, a compiler can represent an application, and/or a graph, using high level language (HL), intermediate level (IL), and/or low level (LL) language constructs and/or statements that can represent operators, operands/results of operators, and/or interconnections of the nodes and/or allocation of CGR hardware to execute the application. HL, IL, and/or LL representations can be, or can represent, an application graph or model. HL, IL, and LL language

constructs/statements can describe nodes and edges of a graph, and/or instructions for executing the graph (i.e., executing the application as represented by the graph) on CGR hardware. HL, IL, and/or LL language constructs and/or statements can include compiler generated mapping alternatives and/or decisions as to how to map the application to CGR hardware for execution.

(84) A compiler can generate a high level graph representation (“HLR”) of an application. The compiler can utilize an HLR, for example, to analyze overall execution elements of the application, and/or to determine initial alternatives for mapping operations of the application to CGR hardware, such as tiling, section cut, and/or parallelization factors in mapping the application.

(85) A compiler can generate, for example, an IL representation (ILR) of the graph that can incorporate mapping alternatives and/or decisions. For example, a compiler can translate an HL graph into an ILR such as an AIR graph and/or a TLIR graph. A compiler can compile, or translate, an ILR to an LL representation (LLR), such as a RAIL representation, that can describe configuration and/or execution instructions to execute the application using particular CGR hardware and/or configurations. The LLR can be suitable for generating application execution code specific to the CGR hardware, such as a PEF and/or configuration files. An ILR and/or LLR can be textual and or graphical, and can be another form of an application, or subset of an application.

(86) A compiler can analyze graphs to determine execution parameters corresponding to CGR hardware allocated to execute the application. For example, a compiler can analyze an ILR (e.g., AIR) or LLR (e.g., RAIL) to determine execution latencies, processor/memory utilizations, and various other such metrics of application execution based on an IL or LL graph that includes CGR hardware resource allocations and/or execution on CGR hardware.

(87) FIG. 4B illustrates example MAC **420**, which can provide functions of a MAC such as MAC **416** in FIG. 4A. FIG. 4B depicts MAC **420** comprising MAC front end **422**, HL optimizer **424**, mapper **426**, IR out **430**, and estimator **428**. In implementations, MAC front end **422** can comprise, for example, an API to input an application and/or application programming statements to compile for execution by a CGRS, shown in FIG. 4B as app **440**. MAC front end **422** can comprise interfaces and/or functions to access hardware descriptions of the CGRS, to access or interact with other components of a compiler that includes MAC **420**, and/or to access or interact with components of a host processor and/or the CGRS. MAC front end **422** can convert an application or application, such as app **440**, to a graph and/or an intermediate representation (IR), for MAC **420** to determine mapping decisions to execute app **440**.

(88) HL optimizer **424** can perform high level optimization of app **440** and/or a graph of app **440**, such as fusing operators (nodes) of a graph into higher level operators, eliminating no-ops and/or redundancies within app **440**, and/or compute derivatives (e.g., Autodiff). In implementations, a compiler can determine a mapping of an application (e.g., operators and tensors included in a graph of an application) to CGR hardware resources for execution of the application. Mapper **426** can be a mapper component or function of MAC **420** that can determine mapping decisions to include in a mapping decision space, such as tiling, section cut, and/or parallelization decisions for mapping app **440** to CGR hardware for executing app **440**.

(89) Mapper **426** can utilize estimator **428** to determine, for example, model execution metrics such as computational latencies of CGRPs executing operators of app **440**, data transfer latencies among memories of CGR hardware (e.g., memories of CGRPs executing operators of app **440**), computational throughput among CGRPs executing operators of app **440**, and/or amounts of memory required for operands/results tensor data of operators of app **440**. Mapper **426** can output mapping decisions to IR out **430** and IR out **430** can translate, or otherwise convert, the mapping decisions to an intermediate representation of app **440** that includes mapping decisions to execute app **440** on the CGR hardware.

(90) As pipelining operations of a dataflow application is an essential aspect of executing the application on CGR hardware, FIG. 5 illustrates an example portion of a graph that can form pipelines among operators in the graph. In FIG. 5, graph **500** is shown comprising operator nodes

N1, N2, N3, N4, and N5 connected by directed edges shown as arrows from one node to another in FIG. 5. If an output tensor of one operator in a graph can share a dimension with an input tensor of another operator that takes the output tensor data, the two operators can potentially form a pipeline based on that shared dimension.

(91) FIG. 5 illustrates example pipelines of nodes N1-N5 based on shared dimensions of the nodes' input and output tensors. A compiler can analyze a graph to identify dimensions on which successive operators of the graph can parallelize their computations in a pipeline. The compiler can, for example, associate dimensions of input/output tensors with a dimension name (or, "Named DIM") as a means to identify tensor dimensions for determining shared dimensions. Tensors having dimensions with the same Named DIM can potentially form pipelines based on the shared dimension corresponding to that Named DIM.

(92) In FIG. 5, suppose that all of nodes N1-N4 have input and output tensors having multiple (e.g., 3) dimensions. Where a dimension of an output tensor of one node and a dimension of an input tensor of another (successor) node share the same Named DIM (that is, share on dimension, among the dimensions of their respective input/output tensors) the nodes (operators) can perform computations in parallel to form a pipeline. In FIG. 5, pipeline 502 represents a pipeline comprising nodes N1-N4, but not N5. In this example, nodes N1-N4 can form pipeline 502 based on a shared dimension (e.g., a dimension having the same Named DIM, say DIM "A") among the dimensions of their respective output/input tensors. However, in this example N5 can have input tensors that do not share DIM A, or any other dimension, with tensors of node N4, such that N5 cannot be included in pipeline 502 or any pipeline based on DIM A.

(93) Nodes of a pipeline can form nested pipelines (pipelines within another pipeline) based on different dimensions among their output/input tensors. As illustrated by the example of FIG. 5, pipeline 502 can comprise nested pipeline 504, and pipeline 504 can comprise nested pipeline 506. Each of pipelines 502, 504, and 506 can be pipelines based on shared tensor dimensions different from that of other pipelines. For example, while pipeline 502 can be a pipeline formed based on shared dimension DIM A, pipeline 504 can be a pipeline formed based on DIM "B", which can be shared among tensors of nodes N2, N3, and N4 but not tensors of nodes N1 and N5. Pipeline 506 is shown comprising nodes N2 and N3, which can be a pipeline formed on dimension DIM "C" shared among tensors of nodes N2 and N3 but not shared by tensors of nodes N1 and N4.

(94) While not shown in the example of graph 500, a node can output tensors to multiple other nodes of the graph (e.g., graph 500 can be a subgraph of a larger application graph that includes operator nodes in addition to those shown in graph 500, and nodes of graph 500 can output tensors to those additional nodes). Thus, nodes among nodes N1-N4 can output tensors to other operators not shown explicitly in graph 500; further, nodes N1-N4 can be included in pipelines based on shared dimensions of tensors of those other nodes.

(95) A "scope" of a pipeline can correspond to the set of operator nodes that can form the pipeline. For example, in FIG. 5 pipeline 502 has a scope comprising operator nodes N1-N4, pipeline 504 has a scope comprising operator nodes N2-N4, and pipeline 506 has a scope comprising operator nodes N2 and N3. However, as a node can be included in only one pipeline at any time for executing the operators in a pipeline, two pipelines cannot have the same scope to execute the operators.

(96) In implementations, a MAC can analyze an application (e.g., a graph of the model) to determine mapping factors included a mapping decision space, such as mapping factors in decision space 400 of FIG. 4A. A MAC can analyze an application or graph to determine operators that can form pipelines, and alternative pipelines, and associated sections including the pipelines, and can include the pipelines in a decision space (e.g., among section and HW boundaries 412 of decision space 400 in FIG. 4).

(97) Based on mapping factors of a mapping decision space, a CGRS compiler can determine alternative configurations of nodes of a graph, such as alternative pipelines and/or sections of a

graph. Based on the decisions, among other mapping factors in a mapping decision space, a CGRS compiler can determine alternative mappings of the graph to CGR hardware to execute the application.

(98) To determine mapping alternatives, and/or mapping decisions, a mapper can make multiple passes (analyses) over a graph of an application model. For example, a mapper can make a tiling pass over the graph to partition input (operand) and output (results) tensors of each operator into smaller segments (tiles and/or slices); a sectioning pass to group operators into sections of the graph; and/or a parallelization pass to determine parallelization alternatives of differing tiling and/or sectioning alternative, and their associated PAR factors, degrees of parallelization (parallelization factors) associated with the section and tiling decisions. As a result of each pass over the graph, the mapper can generate an auxiliary graph, and/or modify an input graph, that can reflect tiling and/or sectioning alternatives and/or decisions. Subsequent passes can utilize the auxiliary graph as an input and can modify an input and/or auxiliary graph based on results of those passes.

(99) Typically, a dataflow application program is written in a high level application programming language, such as Python, that can represent application constructs, such as operator functions, operator operands/results, operator topologies, and dataflow through the operators. However, an application language generally seeks to be hardware independent and, thus, does not represent underlying hardware constructs, such as CGR hardware resources. Consequently, a CGRS compiler can generate an application graph or IR of the application that can express CGR system and hardware structures and/or dataflow, and relationships between application program constructs/algorithms and CGR hardware to execute the application program.

(100) A CGRS compiler can translate an application program to a high level representation (HLR) for purposes of determining mapping decisions such as tiling and sectioning, and corresponding high level resource allocation and scheduling decisions. A compiler can analyze an application to determine CGR hardware mapping alternatives and/or elect mapping decisions, such as tiling and/or sectioning of operators, to allocate CGR hardware resources to execute the application, and/or schedule execution of operators on CGR hardware.

(101) As previously described, an application can be represented as an HLR (e.g., a graph), and can utilize HLR constructs and syntax to describe an application. However, an HLR can require further transformation (compilation) to a low level intermediate representation (LLIR) that can express a mapping of the application to particular CGR hardware resources and data flow among these resources. LL constructs and/or syntax (of an LLIR) can represent attributes of CGR hardware such as processor throughput, pipeline stage latencies, memory capacity and/or bandwidth, memory latencies, and/or other performance and/or structural attributes of CGR hardware that can be allocated to an application during execution. An LLIR can closely represent particular CGR hardware resources and dataflow through the resources (e.g., flow through processors, memories, pipeline stages, and/or intermediate memory buffers) as operators of an application execute on the CGR hardware.

(102) FIG. 6 illustrate an example dataflow compilation process, from an HLR of an application to an LLIR comprising particular hardware resources and associated attributes. In FIG. 6, example compiler 600 comprises HLR 602, decisions 608, and LLIR 604. HLR 602 can comprise an application graph representing operators, operands/results, and/or dataflow of an application (not shown in FIG. 6). In implementations compiler 600 can comprise a compiler such as compiler 320 in FIG. 3.

(103) HLR 602 is shown in FIG. 6 comprising operator nodes N1-N6 (collectively, “nodes 602”) of the application connected in a particular dataflow topology originating with producer node N1 and culminating with consumer node N6. Compiler 600 can generate HLR 602 (or, an alternative HLR) from an application or an alternative application model, or can receive HLR 602 as an input.

(104) Compiler 600 can analyze HLR 602 to determine mapping decisions, CGR hardware

resource allocation decisions, and/or application execution scheduling decisions (e.g., alternative decisions and/or elected alternatives), illustrated in FIG. 6 as decisions **608**. Decisions **608** is shown in FIG. 6 including tiling **610**, parallelization **612**, sectioning **614**, allocations **616**, and scheduling **618**. Tiling **610** can include, for example, tiling decisions, and/or alternative decisions, to partition operands and/or results tensors of nodes among nodes **602**. Sectioning **614** can include, for example, sectioning decisions, and/or alternative decisions, to group nodes **602**. Parallelization **612** can include, for example, parallelizing decisions, and/or alternative decisions, to increase execution parallelism among operations of nodes among nodes **602**. While not shown explicitly in FIG. 6, in implementations compiler **600** can include a MAC, such as MAC **420** in FIG. 4B, and the MAC can determine decisions among decisions **608**.

(105) Based on particular decisions among decisions **608** compiler **600** can determine, for example, that nodes N3, N4, and N5 can form a pipeline (or, pipelines) and compiler **600** can group nodes N3, N4, and N5 in section cut **602B**. Based on other decisions among decisions **608** compiler **600** can determine that nodes N2 and N6 can form a pipeline and can group nodes N2 and N6 in section cut **602C**. However, based on those, and/or other decisions among decisions **608**, compiler **600** may determine that node N1 is not able to form a pipeline with either node N3 or node N6, such that based on these decisions compiler **600** can determine to include node N1 by itself in section cut **602A**.

(106) Among decisions **608** determined by compiler **600**, allocations **616** can include, for example, CGR hardware resource allocation decisions, and/or alternative decisions, to execute operations of nodes among nodes **602** based on other decisions, or decision alternatives, such as among tiling **610**, parallelization **612**, and/or sectioning **614**. Scheduling **618** can include, for example, decisions, and/or alternative decisions, to schedule execution of nodes among nodes **602** on CGR hardware resources included in allocations **616**.

(107) Allocations **616** can comprise high level representations of CGR hardware resources. In translating an HLR to an LLIR, a CGRS compiler can map, or allocate, particular instances of CGR hardware to higher level abstractions of those resources, such as resources that can be included in allocations **616** or other mapping/execution decisions such as illustrated by decisions **608**. To illustrate such a mapping, FIG. 6 shows LLIR **604** comprising processing/memory units PU/MU **606A**, PU/MU **606B** and PU/MU **606C** (collectively, “PU/MUs **606**”). PU/MU **606A**, PU/MU **606B** and PU/MU **606C** can comprise, for example, processor and/or memory units of a CGRP allocated to execute nodes among nodes **602** and/or mapped from corresponding section cuts of HLR **602** (e.g., sectioning **614**). While not shown explicitly in FIG. 6, LLIR and/or PU/MUs **606** can include, for example, configurations, identities, operating characteristics, and/or attributes of corresponding CGR hardware resources that can be used to evaluate performance effects of the mappings of HLR **602** to CGR hardware as shown in the example of FIG. 6.

(108) Typical programming language compilers (e.g., for general purpose, or even special purpose, computers/processors) translate high level language program statements to a low level program (e.g., assembler or machine language). However, as can be seen from the example of FIG. 6, a CGRS compiler can compile an application to an HLR (e.g., a graph and/or text utilizing HL constructs/syntax) and/or can analyze an HLR to determine high level mapping/allocation/execution decisions (hereinafter, “high level decisions”). A CGRS compiler can further compile an HLR to an LLIR that can represent execution aspects of particular CGR hardware unit allocation and execution scheduling. As illustrated in the example of FIG. 6 compiler **600** can generate LLIR **604** to incorporate aspects of high level decisions among decisions **608** in lower level graphical or language constructs representing particular CGR hardware resources, such as among PU/MUs **606**.

(109) It is also common for typical programming language compilers, and even other compilers for dataflow applications, to translate high level language program statements to low level processor instructions and/or execution instructions in a single compilation. However, a CGRS compiler can

determine many alternative mapping decisions to execute a given application model, such as represented by HLR **602**. For example, compiler **600** can potentially partition HLR **602** into many alternative tiles and sections, and can determine potentially many corresponding hardware resource mapping and allocation alternatives. Particular high level application execution decisions, such as illustrated by the example of decisions **608**, can perform much differently from each other when compiled for execution to an LLIR.

(110) As just described, however, an HLR may not represent, or well represent, low level attributes of CGR hardware allocated to an application, such that the effects of high level decisions can be evaluated only after compilation of the HLR to an LLIR. In some cases it may not be possible to know that high level decisions expressed in an HLR are actually feasible to execute the application model. For example, it may not be possible to determine from an HLR that particular high level decisions can fit on CGR hardware allocated or available to execute the application, or that the high level decisions can execute the application within particular execution objectives, such as execution runtime and/or CGR hardware utilization. Some high level decision alternatives, determined based on an HLR, can perform better, while others can perform much worse, when evaluated after compilation of an HLR to an LLIR.

(111) Thus, there can be a cyclic dependency between components, or component levels, of a CGRS compiler, in which evaluating application execution effects of high level compilation decisions depend on compilation to an LLIR or other low level expression. Accordingly, it can be advantageous for a CGRS compiler to evaluate alternative high level decisions, based on an HLR, with respect to effects visible in an LLIR result of compiling the HLR, and to determine which decision alternatives can improve application execution performance and/or achieve particular application execution objectives.

(112) To evaluate effects of alternative high level decisions, an “iterative intermediate” compiler can perform “trial” compilations of high level decisions, based on an HLR, to an LLIR. Based on the LLIR of those decisions, the compiler can evaluate how particular high level decisions affect, or how well the decisions achieve, particular application execution objectives in light of CGR hardware that can execute the application. For example, based on an LLIR compilation from an HLR, an iterative intermediate compiler can assess effects of particular high level decisions on a number, size, and/or type of CGR hardware resources (e.g., processors and/or memories) required to execute the application. An iterative intermediate compiler can assess effects of particular high level decisions in achieving an application execution objective, such as application execution latency, CGR hardware utilizations, and/or allocation of particular CGR hardware resources to execute an application.

(113) To improve determine alternative high level decisions, and/or to elect more optimal decisions, an iterative intermediate compiler can feed back, to high level decision components of the compiler (e.g., a MAC or mapper), results of evaluating effects of high level decisions in an LLIR compilation output. The compiler can determine additional or alternative high level decisions, and can include these decisions in the HLR (e.g., modify HLR mapping, allocation, and/or scheduling decisions), based on feedback from trial LLIR compilations.

(114) FIG. 7 illustrates example method **700** for a compiler to perform an iterative compilation of an application to determine and incorporate high level decisions in an HLR of the application; compile elements of the HLR to an LLIR; use the LLIR to evaluate low level execution effects of high level decisions included in the HLR; and, using the low level evaluation results, determine alternative high level decisions. For purposes of illustrating the method, but not intended to limit implementations, method **700** is described as performed by an iterative intermediate type of CGRS compiler (with reference to FIG. 7, “the compiler”). Also to illustrate the method, the compiler is considered to include a high level (HL) compiler (e.g., a MAC, such illustrated by MAC **420** in FIG. 4B) and a low level (LL) compiler that can compile an HLR to an LLIR (e.g., transform AIR constructs and/or statements to RAIL constructs and/or statements). However, this is not intended

to limit implementations and one of ordinary skill in the art will appreciate that method **700**, and/or operations of method **700** or equivalents of operations of method **700**, can be performed by any type of compiler, or components of a compiler, suitable to compile a dataflow application for execution by a dataflow computing system.

(115) Turning to FIG. 7, in operation **702** of method **700** the compiler generates or receives an HLR of an application. In Operation **702** the compiler can generate the HLR from an application, such as a graph, or can receive an HLR generated previously (e.g., in a previous operation or prior compilation of the application). As previously described, a mapping component of a compiler (e.g., a MAC component of the compiler) can determine high level mapping decisions to execute the application, such as such as illustrated in the example of decisions **608** in FIG. 6, to execute the application on CGR hardware. Additionally, or alternatively, in operation **702** the compiler (e.g., MAC of the compiler) can determine high level decisions to execute the application on CGR hardware. The HLR can include, or be associated with such high level decisions. For example, an HLR can comprise an AIR representation, and a MAC can transform mapping decisions to AIR constructs and/or statements. An AIR can include or represent mapping decisions and/or CGR hardware scheduling and parallelism decisions.

(116) Initially, in operation **702**, the compiler can receive or determine initial high level decisions for all of an HLR or for only a subset of the HLR. In subsequent iterations that repeat operation **702**, the compiler can determine or receive alternative high level decisions for the HLR, or for a portion of the HLR, such as for a particular section, or set of sections, of the HLR.

(117) In operation **704**, the compiler selects one or more candidate high level decisions (hereinafter, with reference to method **700**, “candidate decisions”) from among mapping decisions to execute the application. The compiler can select, in operation **704**, all high level mapping decisions included in, or associated with, an HLR or, alternatively, can select one or more particular high level decisions, to compile to a low level representation. In operation **704**, an HL compiler of the compiler, for example, can select the candidate decisions.

(118) Candidate decisions can correspond to, or be represented within, a section cut of an application, such as illustrated in the examples of sections **602A**, **602B**, and **602C** in FIG. 6. A particular section cut can, for example, incorporate particular computations of the application (e.g., operator nodes in a graph), operand/results tiling, hardware allocation, parallelization of operators and/or computations of an application, and/or hardware execution high level decisions. Thus, in operation **704**, in selecting candidate decisions the compiler can select one or more section cuts of the HLR for trial compilation to an LLIR and evaluation of the candidate decisions to execute the application on CGR hardware. In the example of method **700**, it will be appreciated by one of ordinary skill in the art that the compiler can trial compile a complete HLR, one or more section cuts of an HLR, and/or particular high level decisions.

(119) In operation **706**, the LL compiler initiates a trial compilation of the candidate decision(s), selected in operation **704**. Compiling the candidate decisions, in operation **706**, can include transforming the candidate decisions from the HLR (e.g., representations of operators, section cuts, and/or high level decisions in an AIR) to LLIR constructs, such as RAIL constructs and/or statements, corresponding to execution of the application on particular CGR hardware.

Transforming the candidate decision to an LLIR can include transforming the candidate decisions to low level (e.g., RAIL) allocations of particular CGR hardware resources and/or execution scheduling of the CGR hardware resources. Transforming the candidate decisions to an LLIR, in operation **706**, can determine, for example, a number and/or configuration of processors (e.g., CGRPs) and/or memories, and/or sizes of memories, to allocate for executing the candidate decisions. Scheduling decisions included in an LLIR can include, for example, scheduling execution of particular CGRPs to compute particular functions of the application, and/or scheduling transfer of application data and/or computational results among CGRS memories and/or CGRPs.

(120) Compiling the candidate decisions to an LLIR, in operation **706**, can determine whether particular computations of nodes in a section cut can be performed in parallel, can be pipelined, or must be staged, based on particular CGR hardware resources represented in the LLIR for executing the candidate decisions. Compiling the candidate decisions to an LLIR, in operation **706**, can determine memory and processor requirements to execute the section and, based on the memory required and potentially underutilized processors, the compiler can determine whether particular node computations should be recomputed versus checkpointed in a memory.

(121) In operation **708**, the compiler evaluates the LLIR (e.g., CGR hardware resources, and/or configurations of CHR hardware resources, of the LLIR to execute the candidate decisions) to determine a trial result. A trial result can comprise an indication of whether or not the candidate decisions achieve a particular optimization criterion or application execution objective. A trial result can comprise, or be based upon, low level execution metrics (LL metrics) corresponding to the LLIR resulting from the trial compilation in operation **706**.

(122) The compiler can, in operation **708**, compute, or otherwise determine (e.g., from a specification or description of the CGR hardware) low level execution metrics corresponding to the LLIR, such as corresponding to RAIL allocations and/or scheduling of CGR hardware in an LLIR. LL execution metrics can include, for example, a number of processors (e.g., CGRPs) to execute the candidate decisions (e.g., a section cut included in the candidate decisions); an amount of memory used to execute the candidate decisions; a number of parallel computations to execute a section cut; processor latencies to execute the candidate decisions; and/or a utilization of CGR hardware processors to execute the candidate decisions. LL execution metrics can include, for example, a number of stage buffers required to execute the candidate decisions; a size of one or more memories to allocate to execute a section cut; memory transfer latencies to execute the candidate decisions; and/or a utilization of memories, or particular memories (e.g., on-chip and/or off-chip memory) to execute the candidate decisions.

(123) In operation **708** the LL compiler (for example) can determine a trial result and can output (feedback) the trial result to, for example, the HL compiler. In operation **710** the HL compiler can evaluate the trial result to determine if the candidate decisions achieves an application execution objective. An execution objective can include, for example, a target (objective) hardware utilization, a target memory utilization, a target level of execution parallelism, and/or a target execution latency. An execution objective can include, increasing or, alternatively, decreasing utilization of particular CGR hardware, such as particular processors and/or memories. An execution objective can include increasing parallelism among CGR hardware executing the application, and/or among operators of the application. An execution objective can include decreasing execution latency of the application, or of particular computations of the application, or decreasing a number of stage buffers and/or transfers of data among stage buffer and/or memories, in executing the application.

(124) Evaluating the trial result in operation **710**, can include determining how closely the trial compilation of the candidate decisions achieves or, alternatively, by how much a trial compilation of the candidate decisions fails to achieve, a particular execution objective (e.g., a comparative number of processors to allocate, a comparative amount of memory utilized, or a comparative execution latency, associated with the candidate decisions).

(125) An execution objective can include, or can be associated with, for example, threshold values of corresponding LL execution metrics. A threshold value can comprise, for example, a value that, if exceeded (or, alternatively, not exceeded) by an LL execution metric, can indicate that an associated high level decision (or, decisions) is/are (or, alternatively, is/are not) optimal to execute the application. An LL execution metric can have an upper and lower threshold value, such that the LL execution metric being within (or, alternatively, being not within) a range of values between the upper and lower threshold values can indicate that the candidate decisions achieve or, alternatively, do not achieve an execution objective.

(126) In operation **710** the HL compiler can compare the trial result from operation **708** to trial results from trial compilations of other high level decisions included in an HLR. If a trial result for the candidate decisions compares more favorably to an execution objective than a trial result of an alternative high level decision, the HL compiler can determine to implement the candidate decisions for executing the application. For example, if the trial result from operation **708** increases a number of processors, or reduces and amount of memory used or execution latency, to execute the application, in operation **710** the HL compiler can determine to implement the candidate decisions in lieu of one or more alternative high level decisions.

(127) Based on the trial result from operation **708**, the HL compiler can determine alternative high level decisions to execute the application. An alternative high level decision can comprise a new high level decision not previously included in the high level decisions, and/or can comprise a modification of the candidate decisions selected in operation **704**. For example, if the trial result from operation **708** does not achieve an application objective, the HL compiler can determine one or more alternative mapping decisions that may, potentially, compare more favorably to (e.g., improve a mapping decision in comparison to) the execution objective. The HL compiler can include the alternative mapping decisions in mapping decisions for additional trial compilations and/or selected by the compiler or an application developer (e.g., in operation **714**), to execute the application.

(128) In operation **712**, the compiler determines whether or not to perform additional trial compilations. In operation **712** the compiler can determine to perform additional trial compilations, for example, based on the trial result from operation **708**, for the candidate decisions, not meeting an execution objective, and/or based on the HLR including additional high level decisions to evaluate by trial compilation (e.g., the HLR including additional high level decisions determined by the HL compiler based on the trial result from operation **708**).

(129) Alternatively, in operation **712**, the **712** the compiler can determine to not perform additional trial compilations, for example, based on the trial result determined in operation **708** meeting an execution objective, or that there are no additional high level decisions to evaluate. In operation **712** the compiler can determine to not perform additional trial compilations, for example, based on a limit to trial compilation execution time or resources to perform trial compilations.

(130) If the compiler determines in operation **712** to perform additional trial compilations, the compiler can repeat operations **702-712**. Alternatively, if the compiler determines, in operation **712** to end trial compilations, in operation **714** the compiler ends the trial compilations. In operation **714** ending the trial compilations can include the compiler selecting one or more of the candidate decisions compiled and evaluated, in operations **704-708**, to execute the application.

(131) In operation **714** the compiler can output trial results determined in operation **708**, and/or associated high level decisions, to a developer of the application (e.g., in a file of mapping decisions and/or trial compilation results). Based on the trial results output, developer can select from among the high level decisions to execute the application (e.g., include directives to the compiler for compiling the application based on the selected high level decisions).

(132) In ending the trial compilations in operation **714**, the compiler can output and/or store trial results determined in operation **708**, and/or high level decisions selected by the compiler to execute the application. For example, the compiler can output and/or store LLIRs produced by trial compilation in operation **706**, trial results determined in operation **708**, and/or compiler metrics corresponding to the compiler performing the trial compilations (e.g., runtimes, computational resource consumption, etc. associated with one or more operations of method **700**). The compiler can output the metrics and/or results to, and/or store the metrics and/or results in, a memory and/or a storage medium, device, or subsystem. In operation **714** the compiler can output the metrics and/or results to a network and/or another computer or computing system communicatively coupled to the computer or computing system executing the compiler.

(133) FIG. **8** illustrates an example system for performing iterative trial compilations of an

application program to execute on a dataflow system. Example computing system **800** can perform a method similar or equivalent to method **700**, of FIG. 7, and/or operations thereof. In FIG. 8, computing system **800** is shown comprising APP **802**, processor **804**, and compiler **808**. In implementations, APP **802** can comprise any type of application for execution by a dataflow computing system, such as dataflow and data parallel applications previously described.

(134) In implementations, compiler **808** can comprise an iterative intermediate compiler and can include programs and/or logic circuits to compile APP **802** for execution by a dataflow computing system, such as a CGRS. Compiler **808** can compile APP **802** to execute, for example, on CGR hardware of a CGRS. Compiler **808** can, for example, generate an HLR based on APP **802**, shown in FIG. 8 as HLR **810**, or can receive HLR **810** as an input. HLR **810** can comprise, for example, a graph and/or HL statements describing computational operators and/or functions and/or operands/results of APP **802**. HLR **810** can describe a topology of operators, and/or dataflow among the operators, of APP **802**. HLR **810** can comprise AIR constructs and/or statements, for example.

(135) Processor **804** can comprise a processor, or processors, of a computing system that can execute programs of a compiler such as compiler **808**. Processor **804** can comprise any type of computing processor, or combination of processors, capable of executing programs of a compiler. Compiler **808** is shown coupled to processor **804** by means of interface **806**, which can comprise any type of hardware and/or software interface to execute programs of compiler **808** on processor **804**. In implementations interface **806** can be, for example, an API for executing compiler **808** on processor **804**, for compiler **808** to receive input data, such as APP **802** and/or HW SPEC **812**, and/or for compiler **808** to output data, such as an IR, of APP **802**, high level decisions to execute APP **802**, and/or results of trial compilations of APP **802** or portions of APP **802**. Interface **806** can include functions, programs, and/or interfaces of an operating system (not shown in FIG. 8) of computing system **800** and the operating system can facilitate or control execution of programs of compiler **808**.

(136) Compiler **808** can include a high level compilation component, shown in FIG. 8 as HL compiler **820**. HL compiler **820** can comprise, for example, an algebraic graph compiler, such as algebraic graph compiler **322** in FIG. 3, or a MAC component or function of compiler **808**, such as MAC **420** in FIG. 4B. HL compiler **820** is shown in FIG. 8 comprising high level decision components (or, functions) tiling **832**, parallelizing **834**, sectioning **836**, and scheduling **838**, which can individually and/or in combination determine high level decisions to execute APP **802** on a dataflow computing system, or elements of APP **802** (nodes and/or operand/results tensors of a graph of APP **802**), to CGR hardware to execute APP **802**. In FIG. 8, HL decisions **822** can be an output of HL compiler **820** (and/or high level decision components of HL compiler **820**), such and can include high level decisions such as high level decisions illustrated by decisions **608** of FIG. 6 to execute APP **802**.

(137) Compiler **808** can include, receive, or can have access to (e.g., in a memory or an input file), a CGR hardware description, shown in FIG. 8 as HW SPEC **812**. HW SPEC **812** can include a detailed specification of hardware of a computing system (e.g., CGR hardware of a CGRS) for executing APP **802**. HW SPEC **812** can include a description of hardware configurations and/or availability of particular hardware (e.g., processors and/or memories) to execute APP **802**. HL compiler **820** can utilize HLR **810** and HW SPEC **812** to determine high level mapping decisions, among HL decisions **822**, to execute APP **802** on hardware described in HW SPEC **812**.

(138) Compiler **808** can include a low level compilation component and/or function, shown in FIG. 8 as LL compiler **824**. LL compiler **824** can comprise a component and/or function of compiler **808** that can use hardware specifications in HW SPEC **812**, and high level decisions included in HL decisions **822**, to compile HLR **810** to an LLIR, shown in FIG. 8 as LLIR **826**. LLIR **826** can comprise any form of low level representation such as previously described. In particular, LLIR **826** can include constructs and/or language statements representing or corresponding to particular

hardware included in HW SPEC **812** and/or representing or corresponding to execution flow of APP **802** (e.g., a dataflow, pipeline, and/or parallel execution of operations of APP **802**). LLIR **826** can comprise RAIL constructs and/or statements, for example.

(139) Compiler **808** can perform trial compilations, such as illustrated in the example of method **700** of FIG. 7, of HLR **810** and decisions among HL decisions **822**. HL compiler can perform operations of operations **702** and/or **704** of example method **700** in FIG. 7 to determine high level decisions in HL decisions **822**. LL compiler **824** can, for example, perform operations of operation **706** of example method **700** in FIG. 7 to compile HLR **810** and HL decisions **822** to LLIR **826**. LL compiler **824** can perform trial compilations of all of HLR **810** utilizing alternative decisions among HL decisions **822**. LL compiler **824** can perform trial compilations of all elements of HLR **810** utilizing alternative or particular decisions among HL decisions **822**. Additionally, or alternatively, LL compiler **824** can perform trial compilations of subsets of elements of HLR **810** (e.g., as illustrated in the example of method **700** in FIG. 7, a trial compilation of a section of HLR **810**) utilizing alternative or particular decisions among HL decisions **822**.

(140) LL compiler **824** (or, a compilation evaluation component of Compiler **808** or LL compiler **824**, not shown explicitly in FIG. 8) can evaluate results of a trial compilation of HLR **810** and decisions among HL decisions **822**. For example, LL compiler **824** can perform operations of operation **708** of example method **700** in FIG. 7 to evaluate results of a trial compilation of HLR **810** and decisions among HL decisions **822**. FIG. 8 illustrates computing system **800** including objectives **814**, which can comprise application execution objectives, such as previously described in reference to method **700** of FIG. 7. Objectives **814** can correspond, for example, to LL execution metrics and/or threshold values of LL execution metrics. To evaluate results of a trial compilation, LL compiler can determine LL execution metrics resulting from a trial compilation and can compare the execution metrics to execution objectives included in objectives **814**.

(141) LL compiler **824** can output results of a trial compilation, and/or evaluation results of a trial compilation, to trial results **828**. HL compiler **820** can receive (e.g., from LL compiler **824**), or can otherwise access (e.g., in a memory or file of a storage medium), trial results **828** corresponding to a trial compilation by LL compiler **824** of some or all of HLR **810** and some or all of decisions among HL decisions **822**.

(142) HL compiler **820** can evaluate results of trial compilations, performed by LL compiler **824**, to determine and/or select (or, alternatively, discard) decisions among HL decisions **822** to execute APP **802**, and/or to determine alternative high level decisions among HL decisions **822**. HL compiler **820** can, for example, perform operations of operation **714** of method **700**, in FIG. 7, to select or, alternatively, to discard decisions among HL decisions **822** to execute APP **802**, and/or to determine alternative high level decisions among HL decisions **822**.

(143) In combination, HL compiler **820** and LL compiler **824** can iteratively compile alternative high level decisions, and/or elements of HLR **810**, to map operations of APP **802** to hardware of a computing system to execute APP **802**, such as CGR hardware of a CGRS to execute APP **802**. HL compiler **820** and LL compiler **824** can perform such iterative, trial compilations and evaluations of trial results to determine more, and/or less, optimal mappings (e.g., decisions among HL decisions **822**) of APP **802** to hardware to execute APP **802**.

(144) Compiler **808**, HL compiler **820**, and LL compiler **824** can comprise software elements (programs), hardware elements (e.g., special purpose compilation and/or graph processors or circuits), or a combination of software and hardware elements. While not shown explicitly in FIG. 8, computing system **800** can include memories, storage media, and/or network interfaces. Programming instructions for Compiler **808**, HL compiler **820**, and/or LL compiler **824** to perform methods and/or operations of the disclosure, can be stored in a storage medium and/or memory included in (but, not shown explicitly in FIG. 8), or communicatively coupled to, computing system **800**, and/or accessible to processor **804**.

(145) In implementations, a storage medium can comprise any form of persistent memory or

medium, such as a disk or optical storage device, or a storage medium of a computer storage subsystem. A memory can comprise any computer memory that can be accessible to compiler **808**, and/or processor **804** while executing programs of compiler **808**. For example, while not shown explicitly in FIG. **8**, computing system **800** can include a main memory, on-chip and/or off-chip memories of a CGRS, cache memories, and/or memories of a CGRS (e.g., PMUs, DRAMs, and/or SRAMs). Network interfaces can include programs and/or hardware of, or to access, a network, such as an Ethernet network and/or the Internet.

(146) Compiler **808** can access input information (e.g., APP **802**, HLR **810** HW SPEC **812**, and/or objectives **814**) in the memories, storage media, and/or via the network interfaces. Interface **806**, for example, can include software (e.g., APIs and/or command line interfaces) and/or hardware interfaces (e.g., buses, I/O links, and/or network interfaces) to enable compiler **808** and/or processor **804** to access the memories, storage media, and/or a network. Compiler **808** can store outputs of Compiler **808** (e.g., HLR **810** and/or modifications to HLR **810**, HL decisions **822**, LLIR **826**, and/or trial results **828**) in the memories, storage media, and/or via the network interfaces. The memories, storage media, and/or network interfaces can store or provide access to programs of Compiler **808**.

(147) The example of FIG. **8** is intended only to illustrate example components, and arrangements thereof, of a system capable of performing iterative compilations of an HLR corresponding to an application model. However, the examples of FIG. **8** are not intended to limit implementations. It would be apparent to one of ordinary skill in the art that such a system can comprise a variety of components, and/or arrangements of components, alternative to the example of FIG. **8**, that such alternative components/arrangements can perform iterative compilation methods and/or operations of the disclosure, and that such alternative implementations will be within the scope and spirit of the disclosure.

(148) As previously described, an application program can be written in a high level application language, such as Python, that generally seeks to be hardware independent and, thus, does not represent underlying hardware constructs, such as particular CGR hardware resources.

Consequently, a dataflow application program (hereinafter, for brevity, “dataflow program”), for example, can have program execution inefficiencies associated with executing the program on actual CGR hardware of a CGRS. Program execution inefficiencies (hereinafter, for brevity, “inefficiency” and “inefficiencies”) in a dataflow program can be categorized, for example, as operator inefficiencies, operator and/or hardware pipeline imbalances, memory stalls, transient overheads, unused hardware resources, and/or underutilized hardware resources.

(149) Operator inefficiencies can correspond to, or depend on, CGR hardware design and function/implementation of low-level Operator Programs (OPPs) to execute operators, and/or execute a set of operators forming a pipeline. OPPs can execute operators of an application (dataflow computations) and can comprise low level processor and/or CGR hardware configuration/execution instructions to implement the operators. OPPs can comprise, for example, templates such as described in reference to the examples of FIG. **3**. A CGRP (e.g., a PCU) can be programmed to execute a particular OPP. Operators in an application can form pipelines and CGR hardware can be configured (e.g., by programming switches of a compute grid, such as a PCU/PMU array) to execute a set of OPPs in parallel as a pipeline.

(150) Operator inefficiencies can, for example, correspond to operator and/or CGR hardware pipeline initialization (warmup) and cooldown (terminating the pipeline); draining (completing remaining) processor computations; and/or, inefficiencies in low level programming of OPPs, and/or sequences of OPPs. A difference between actual execution latency of an operator (or, a set of operators, such as operators of a pipeline) and a theoretical minimum execution latency (sometimes referred to as a “roofline” latency), for example, can indicate an operator inefficiency.

(151) Pipeline imbalances refer to differing execution latencies between CGRPs executing OPPs of differing operators in a pipeline, in which one stage (e.g., one PCU, or set of PCUs) of a pipeline

takes longer or less time than a producer and/or consumer stage, such that some stages may have to suspend processing pending completion of a longer executing stage. Imbalances can result based on tiling alternatives, section cuts, and/or PAR factors associated with section cuts.

(152) Memory stalls refer to circumstances in which a CGRP's memory cannot operate at the same speed (e.g., frequency/bandwidth/data rate) as the processor's computational speed. This can cause a processor to stall (suspend processing) to await completion of memory accesses. Memory stalls can depend on a type of memory used by an OPP, such as on-chip (local to the processor) versus off-chip (remote from the processor as implemented in a different IC). Memory stalls can depend on a type or technology of memory (e.g., SRAM, register stack, or DRAM) used in executing an OPP and/or an OPP or processors pipeline.

(153) Transient effects refers to events that can occur only once, or intermittently, in execution of OPPs in a pipeline. Transients effects can be associated with, for example: pipeline startup (internal to processors, as well as among processors executing operator pipelines of the application; programming reconfigurable CGR hardware (e.g., an RDU) to execute a particular OPP; memory initialization (e.g., fetching initial operands); memory-memory transfers between OPPs and/or internal pipeline stages of a processor; CGRA execution dependencies on, or interactions with, dataflow system operating and/or runtime systems (e.g., receiving input/communicating output from CGRA execution to OS programs, device drivers, runtime execution management, etc.); and/or transients effects that can be determined from steady-state operations of application execution.

(154) Program execution inefficiencies can lead to suboptimal performance in executing a dataflow program on a dataflow system, such as a CGRA based system. Thus, a dataflow application developer (e.g., a dataflow application programmer, or an automated dataflow application generator program or system) can benefit from an analysis of the dataflow program that can identify potential underlying inefficiencies associated with a dataflow program. Thus, in implementations a dataflow computing system, such as a CGRS, can include an “efficiency analyzer” (hereinafter, for brevity, “analyzer”) that can analyze a dataflow program to predict CGR hardware latency and/or utilizations and, based on the predictions, can identify potential inefficiencies in executing the program, or portions of the program, on particular CGR hardware.

(155) The analyzer can determine inefficiency metrics associated with an inefficiency category. In implementations inefficiency metrics can include, for example, a total number of computational cycles (e.g., Floating Point Operations Per Second, or “FLOPs”) wasted (e.g., idle or suspended instruction cycles) due to an inefficiency, and/or an aggregate total of all wasted cycles within a particular inefficiency category. In another example, an inefficiency metric can correspond to a total number of unused processors multiplied by the amount of time for the processors are unused, which can represent a fraction of total program peak cycle rate (e.g., peak FLOPs), and/or a fraction of total execution latency.

(156) An analyzer can determine and organize potential inefficiencies according to an inefficiency category. Inefficiency categories can correspond, for example, to a time-average of sources of underutilized CGR hardware resources during application (or, partial application) execution, and/or other inefficiencies. An analyzer can represent, and/or output, inefficiencies in a particular dataflow program in various formats, and/or organized by inefficiency categories, for an application developer to utilize as a means to revise the program, such as to reduce, and/or eliminate, the inefficiencies.

(157) FIG. 9 illustrates example method 900 to predict hardware utilization during execution of a dataflow program and, based on the predictions, identify and output program execution inefficiencies. For purposes of illustrating the method, the method is described as performed by an application analyzer (in further reference to method 900, “an analyzer” and “the analyzer”) to analyze a dataflow program. In implementations, the analyzer can be a standalone component of a dataflow systems, and/or can be included in, or integrated with, a dataflow program compiler.

However, this is not intended to limit implementations and it would be appreciated by one of ordinary skill in the art that the examples of method **900** can apply to components of a computing system other than an analyzer component such as previously described, and/or that the examples and/or operations of method **900** can apply to application programs alternative to dataflow application programs.

(158) An analyzer can evaluate an IR (e.g., an HLR and/or LLIR) of a dataflow program to predict CGR hardware utilization and/or identify potential execution inefficiencies and method **900** is further described as the analyzer analyzing an IR. The terms “an IR” and, “the IR”, for purposes of describing method **900**, refer interchangeably to an HLR (e.g., a graph), a LLIR (e.g., a textual representation), and a combination of both. In a dataflow program operators can be included in a logical pipeline. The operators can be implementable as hardware pipelines and the term “meta-pipe” is used herein to refer to a set of operator of a dataflow program that can form a pipeline in the program and can be implementable as hardware pipelines in a dataflow computing system.

(159) Additionally, in the description of method **900**, the term “meta-pipe” is interchangeable to refer to a pipeline of application operators/operand/computational results as a whole (which can be a subset of an application graph), and to a nested pipeline within a pipeline. However, this is also for purposes of illustrating the disclosure and not intended to limit implementations.

(160) An IR can include a mapping of application operators, operands, computational results, and/or dataflow to CGR hardware of a CGRS for executing the application. The IR (e.g., a mapping included in an IR) can include (or, be accompanied by) mapping decisions made by a compiler (e.g., an HL compiler) to map a dataflow program to CGR hardware for execution. The mapping decisions can be included in a decision space or search space of the compiler, and/or can be included in an HL and/or LL representation of a compilation result.

(161) In operation **902** the analyzer selects a set of operators of the dataflow program to analyze for program execution inefficiencies. In operation **902** the analyzer can select the set of operators from an IR, and/or a mapping included in or associated with an IR. In an IR a meta-pipe, for example, can represent dataflow program operators, operands, computational results, and/or dataflow (e.g., a flow of operands and/or computational results among hardware elements executing a dataflow program) and the IR can include a mapping of the meta-pipe for execution on CGR hardware. Accordingly, to further illustrate the method but not intended to limit implementations, in operation **902** the analyzer can select, from an IR, a candidate meta-pipe to analyze for potential application execution inefficiencies. However, this is not intended to limit implementations, and in alternative implementations, in operation **902** the analyzer can, for example, select from the IR a set of operators other than operators included in a meta-pipe, and can perform operations of method **900** using this alternative set of operators.

(162) In operation **904**, the analyzer predicts execution performance of the candidate meta-pipe, and/or particular operators (e.g., low level OPP s) of the candidate meta-pipe, on particular CGR hardware and/or particular types of CGR hardware. To predict the execution performance the analyzer can utilize a hardware description of the CGR hardware to compute predicted execution metrics, such as predicted latencies and/or utilization of CGR hardware, corresponding to executing the candidate meta-pipe and/or operators selected in operation **902**, and/or dataflow of the candidate meta-pipe and/or operators, on particular CGR hardware.

(163) The hardware description can specify, for example, configuration (e.g., type and/or number, interconnections, and/or configurations of hardware elements) and/or operating parameters (e.g., processor and/or data transfer latency, and/or processor computational throughput or instruction cycle rates) of CGR hardware elements. The analyzer can compute the predicted execution metrics based on mapping decisions of an IR to execute operators of the meta-pipe on particular CGR hardware, CGR hardware types, and/or the hardware description.

(164) Based on the predicted execution metrics, in operation **904** the analyzer further determines inefficiency metrics, such as previously described. In operation **906**, based on the inefficiency

metrics the analyzer determines inefficiencies in the dataflow program, such as program execution inefficiencies previously described. In operation **904** the analyzer can compute, for example, an inefficiency metric comprising a difference between aa predicted (e.g., actual) execution latency of an operator or, a set of operators, such as operators of a pipeline, and a theoretical minimum execution latency to execute the operator or set of operators.

(165) In operation **908** the analyzer records inefficiencies determined in operation **906**, and/or predicted execution and/or inefficiency metrics determined in operation **904**, associated with executing the candidate meta-pipe (or, other operators) selected in operation **902**. In operation **908** the analyzer can associate, and/or record, the inefficiencies and corresponding predicted execution and/or inefficiency metrics with inefficiency categories such as previously described.

(166) The analyzer can record the candidate meta-pipe inefficiencies and/or predicted execution and/or inefficiency metrics and/or in a memory, a storage device, and/or in a data structure of a file, for example. The analyzer can include the inefficiencies and/or predicted execution and/or inefficiency metrics in (e.g., as an output to) a decision space, search space, and and/or IR (e.g., an HLR and/or LLIR) representations of a dataflow program. The analyzer can include the inefficiencies and/or predicted execution and/or inefficiency metrics as attributes of operators of the dataflow program.

(167) In operation **910** the analyzer determines if the candidate meta-pipe, selected in operation **902**, includes a nested meta-pipe (inner meta-pipe) that has not been evaluated. If so, in operation **912** the analyzer selects a nested meta-pipe of the candidate meta-pipe and repeats operations **904-910** with the selected nested meta-pipe as a new candidate meta-pipe.

(168) If, in operation **912**, the analyzer determines that the candidate meta-pipe, selected in operation **902**, does not have nested meta-pipes, in operation **914** the analyzer determines if there are other, parent meta-pipes to analyze. If so, the analyzer repeats operations **902-914** to select and evaluate a new parent meta-pipe. If, in operation **912**, the analyzer determines that there are no parent meta-pipes to analyze, in operation **916**, the analyzer outputs the results of the inefficiency analysis.

(169) In operation **916**, the analyzer outputs results of the inefficiency analysis (hereinafter, “inefficiency results”) performed in operations **902-914**. The inefficiency results output in operation **916** can include predicted execution and/or inefficiency metrics determined in operation **904** and/or inefficiencies determined in operation **906**. The inefficiency results can comprise results for all meta-pipes (or, other operators) selected in operation **902** or, alternatively, can comprise results for a subset of meta-pipes (or, other operators) selected in operation **902**.

(170) The inefficiency results output in operation **916** can include suggestions and/or recommendations to modify a dataflow program, and/or mapping decisions, to improve on, or eliminate, inefficiencies included in the results. The suggestions and/or recommendations can correspond to particular design characteristics of the CGR hardware.

(171) The inefficiency results output in operation **916** can comprise, for example, a list of inefficiency data, tabular inefficiency data and/or graphical inefficiency data, or can be organized within a data structure and/or database. In operation **916**, the analyzer can output the inefficiency results to an interface of a compiler or of a dataflow system, such as an API, CLU, and/or GUI of a compiler dataflow system. The analyzer can output the inefficiency results to a memory and/or a storage system.

(172) In implementations an analyzer can, in operation **916** and/or other operations of a method such as **900** in FIG. **9**, output the inefficiency results, or otherwise make the inefficiency results available (e.g., in a file for later use), to an “analysis assistant”. An analysis assistant can present inefficiency results to a dataflow program developer (hereinafter, for brevity, “developer”) in a manner that can assist the developer in interpreting the inefficiency results, and/or to revise a dataflow program to remove or reduce inefficiencies in executing the program on a particular CGRS, and/or particular CGR hardware of a CGRS.

(173) A developer can be a human programmer of a dataflow program; additionally, or alternatively, a developer comprise, or can include, an automated system, such as another computing application functioning to develop, or assist in developing, a dataflow program. Thus a developer can be a human programmer, a computing system designed to program and/or revise dataflow applications, or a combination of a human and a computing system for developing/revising a dataflow program.

(174) An analysis assistant can comprise, for example, a programming component of a compiler and/or a component of a CGRS (or, a computing system coupled to the analyzer, compiler, and/or CGRS). In another example, an analysis assistant can comprise a component of a dataflow compiler of a CGRS and/or can be a component of a computing system communicatively coupled to a computing system, memory, and/or storage medium, having a record of inefficiency results. An analysis assistant can comprise a component of a software development kit for implementing a dataflow program for execution on a particular CGRS.

(175) FIG. 10 illustrates example method **1000** that an analysis assistant can perform to present inefficiency results to a developer. For purposes of illustrating the method, the method is described as performed by an analysis assistant (“an assistant” and “the assistant”, in reference to method **1000**) of a compiler of a CGRS, in presenting inefficiencies within a dataflow program to execute on the CGRS. However, this is not intended to limit implementations and it would be appreciated by one of ordinary skill in the art that the examples of method **1000** can apply to a computing system other than a CGRS and/or can apply to a presentation component of a computing system other than as included in a compiler. It would be further appreciated by one of ordinary skill in the art that method **1000**, and examples and/or operations of the method, can apply to application programs alternative to dataflow application programs.

(176) In operation **1002** the assistant initiates presentation of inefficiency results. The inefficiency results can comprise results such as can be output (e.g., in operation **916**) by an analyzer performing a method such as method **900** in FIG. 9. Inefficiency results can include predicted execution and/or inefficiency metrics, and the predicted execution and/or inefficiency metrics can be associated with, or the assistant can associate the predicted execution and/or inefficiency metrics with, inefficiency categories. A presentation can comprise a particular means of outputting inefficiency results in a particular form or format (e.g., a graphical and/or textual format) that can render the inefficiency results most useful for a developer to identify a source of an inefficiency in a dataflow program and/or revise the dataflow program to remove or reduce the inefficiency. A presentation of inefficiency results can utilize an output interface of a computing system to output and/or display inefficiency results composed by the assistant in a format to assist a developer in identifying execution inefficiencies, and/or sources of execution inefficiencies, in a dataflow program.

(177) A presentation can be interactive (e.g., via an API, CLI and/or other form of output and/or user interface) with a developer in programming a dataflow program. An interactive presentation can assist a developer to analyze inefficiency results, such as formatted inefficiency results, and to interact with a compiler to revise a dataflow program to remove or reduce an inefficiency in the program. For example, in the course of programming a dataflow program, in an interactive presentation a developer can receive formatted inefficiency results associated with analyzing (e.g., by an analyzer) a version of a dataflow program; can revise the program to remove or reduce an inefficiency; can recompile the program to an LLIR; can invoke (explicitly or, alternatively, implicitly in recompiling the program); and, can receive new formatted inefficiency results associated with the revisions to the program. In operation **1002** the assistant can present the inefficiency results to, and/or interact with, a developer via a communication or network interface.

(178) A compiler can include, or can invoke, a presentation program for the compiler to present the inefficiency results in different forms. In operation **1002**, the assistant can utilize a software interface of a computing system to initiate presentation of the inefficiency results. A software

interface can comprise an API, CLI, and/or GUI of a computing system executing or, in communication with, the assistant. The assistant can, in operation **1002**, initiate presentation of the inefficiency results, for example, in response to a request (e.g., received via a software or communications interface) of a developer to present the inefficiency results and/or assist the developer in observing, interpreting, and/or acting on the inefficiency results to revise a corresponding dataflow program. The assistant can initiate the presentation, in operation **1002**, in response to an analyzer outputting (e.g., via an API or CLI of an analyzer or compiler) inefficiency results and/or in response to receiving outputting (e.g., via an API or CLI of an analyzer or compiler) inefficiency results output from an analyzer (e.g., in response to receiving the inefficiency results from an analyzer performing operation **916** of method **900**).

(179) In operation **1004** the assistant receives the inefficiency results and composes formatted inefficiency results. In operation **1004** the assistant can receive inefficiency results, for example, as an output of an analyzer, such as an output of operation **916** of method **900** in FIG. 9. The assistant can receive the inefficiency results from, for example, a memory, a storage device (e.g., as a file stored on a storage device), and/or a communication or network interface.

(180) The assistant can, in operation **1004**, compose the formatted inefficiency results to have a form or format that can assist a developer to visualize and/or otherwise interpret the inefficiencies, to determine sources of inefficiencies, determine CGR hardware utilization associated with inefficiencies, and/or determine revisions to a dataflow program to remove or reduce one or more inefficiencies. For example, in operation **1004** the assistant can compose formatted inefficiency results in a tabular format. A tabular format can comprise, for example, one or more tables of predicted execution and/or inefficiency metrics, and/or tables corresponding to inefficiency categories comprising associated predicted execution and/or inefficiency metrics.

(181) In another example, in operation **1004** the assistant can compose formatted inefficiency results in a graphical format, such as one or more graphs, pie charts, and/or bar charts that can graphically represent inefficiency metrics as compared to optimal or theoretical maximum application execution. Such graphs and/or charts can represent, for example, CGR hardware utilization by operators of a dataflow program mapped to that CGR hardware.

(182) The assistant can, in operation **1004**, compose formatted inefficiency results as a “latency waterfall chart”. A latency waterfall chart can, for example, plot one or more operators' execution latency, as determined in an inefficiency analysis, versus a “roofline” theoretical minimum latency for a particular number of CGR hardware compute units (e.g., PCUs and/or other CGR hardware processors) mapped to execute the operators included. A latency waterfall chart can complement other graphical presentations (e.g., graphs, pie charts, and/or bar charts) in assisting a developer to remove and/or reduce inefficiencies in a dataflow program. For example, a waterfall latency chart combined with a bar chart having a predicted latency bar and another bar chart having a roofline bar can illustrate that OPP inefficiency is zero only if the predicted latency bar equals the roofline bar in a latency waterfall chart having both bars spanning the dataflow of a set of operators through CGR hardware.

(183) In another example, in operation **1004** the assistant can compose formatted inefficiency results as a combination of a waterfall latency chart with a bar chart having a predicted latency bar can illustrate that pipeline imbalance in CGR hardware executing an operator pipeline is zero only if all predicted latency bars are equal. Similarly, a combination of a waterfall latency chart with a bar chart having a predicted latency bar and/or a bar chart having a roofline latency bar can illustrate differences in execution latencies, as a dataflow program is programmed, and predicted and/or roofline latencies.

(184) In operation **1006** the assistant presents the formatted inefficiency results for a developer to use in identifying inefficiencies and/or sources of inefficiencies. Presenting the formatted inefficiency results can comprise the assistant outputting the formatted inefficiency results to an output interface of a computing system, such as a file system, display or monitor interface, or

communications interface, or via another interface such as an API, CLI, or other user interface (e.g., a GUI). The output interface can comprise an interface to output the formatted inefficiency results for use by the developer to interpret the inefficiency results and/or improve the dataflow program in association with the inefficiency.

(185) In an interactive presentation, in operation **1006** the assistant can present the inefficiency results to the developer using an interactive interface of a computing system executing the analysis assistant, or communicatively coupled to the analysis assistant. An interactive interface can comprise, for example, an internet connection, a display of a computing system (e.g., a graphical display), and/or other input/output devices (e.g., audio, video, and/or communications interface devices) of a computing system for communicating with a developer. An interactive interface can comprise an API, CLI, and/or GUI for communicating with a developer. In an interactive presentation, the assistant can, in operation **1006**, present inefficiency results as a developer is actively writing a dataflow program, and/or as an analyzer outputs inefficiency results (e.g., in a case in which a developer revises a dataflow program and repeats an inefficiency analysis).

(186) In operation **1008** the assistant, optionally, determines if there are program revision suggestions and/or recommendations to present to the developer that can improve (remove and/or reduce) inefficiencies in the dataflow program. In operation **1008** the assistant can determine that there are suggestions/recommendations based on an analyzer having generated suggestions and/or recommendations, and/or based on an analyzer and/or the assistant comparing the inefficiencies to optimal or theoretical application execution on CGR hardware. The suggestions/recommendations can correspond to inefficiencies associated with particular operators (e.g., particular OPPs) and/or particular design characteristics of CGR hardware. For example, a suggestion or recommendation can comprise modification(s) of one or more OPPs, associated with particular inefficiencies and/or inefficiency categories, for executing one or more operators on particular CGR hardware (e.g., inefficiencies determined by an analyzer).

(187) The assistant can, determine, in operation **1008**, that there are suggestions and/or recommendations in response to the developer requesting suggestions. A suggestion and/or recommendation determined in operation **1008** can comprise a suggestion and/or recommendation for improving all inefficiencies or for improving only a subset of inefficiencies. If the assistant determines, in operation **1008**, that there are suggestions and/or recommendations to present, in operation **1010** the developer presents the suggestions/recommendations using an output interface, such as an API, a CLI, a file, a computer display, or a user interface. The output interface can comprise an interface to output the suggestions/recommendations for use by the developer to improve the dataflow program in association with the inefficiency.

(188) In operation **1010** the assistant can present the suggestions/recommendations as an output for subsequent use by a developer, such as a file stored in a memory and/or storage medium. The assistant can repeat operations **1008** and **1010** to present a series of suggestions/recommendations, and/or to interact with the developer to determine and present program revision suggestions/recommendations.

(189) In an interactive presentation, in operation **1010** the assistant can present suggestions/recommendations to the developer using, for example, an internet connection, a display of a computing system (e.g., a graphical display), and/or other input/output devices (e.g., audio, video, and/or communications interface devices) of a computing system used by, or in communication with, a developer. In an interactive presentation, the assistant can, in operation **1010**, present suggestions/recommendations as a developer is actively writing a dataflow program, and/or as an analyzer outputs inefficiency results (e.g., in a case in which a developer revises a dataflow program and repeats an inefficiency analysis).

(190) Alternatively, if the assistant determines in operation **1008** that there are no suggestions or recommendations to present or, optionally, does not perform operation **1008**, in operation **1012** the assistant determines if the presentation is complete. In operation **1012** the assistant can determine

that the presentation is complete based on having, for example, presented all inefficiency results and/or revision suggestions, or has presented a particular subset of the inefficiency results and/or suggestions/recommendations. In operation **1012** the assistant can determine that the presentation is complete based on an indication, or request, by a developer and/or analyzer to end the presentation.

(191) If the assistant determines, in operation **1012**, that the presentation is not complete, the assistant can repeat operations **1004-1012**. In repeating operations **1004-1012** the assistant can receive and/or retrieve, in operation **1004**, additional or alternative inefficiency results, and/or can retrieve the inefficiency results previously selected in operation **1004**. In repeating operations **1008-1010** the assistant can present previously determined and/or presented suggestions/recommendations and/or can determine and/or present other (e.g., new) suggestions/recommendations.

(192) Alternatively, if the assistant determines, in operation **1012**, that the presentation is complete, in operation **1014** the assistant ends the presentation. In ending the presentation, in operation **1014**, the assistant can store (e.g., in one or more files) information associated with the presentation, such as the inefficiency results, revision suggestions/recommendations, and/or results of interactions with the developer. (e.g., revisions to a dataflow program). The assistant can store the information in a memory and/or a storage device, and/or can communicate the information to a developer, compiler, analyzer, and/or another component of a CGRS. In operation **1012**, in ending the presentation the assistant can initiate the analyzer repeating the analysis of the previous, or an alternative (e.g., a revised), dataflow program.

(193) FIG. **11** illustrates an example system for performing methods and/or operations of the disclosure, such as method **900** and/or **1000**, and/or methods similar or equivalent to methods **900**, **1000**, and/or operations thereof. In FIG. **11**, computing system **1100** is shown comprising APP **1102**, compiler **1108**, developer **1160**, analysis assistant **1150**, and processors **1104A** and **1104B**.

(194) In implementations, APP **1102** can comprise, for example, a dataflow program and/or other application program for performing functions such as machine learning, natural language processing, image, video, and/audio processing. Developer **1160** can be a human and/or automated (e.g., a by a computing application, such as an AI application for developing applications) application developer, or programmer, to develop or write APP **1102**. Analysis assistant **1150** can interact with developer **1160** to identify and/or correct inefficiencies in APP **1102**.

(195) Processors **1104A** and/or **1104B** can be any type and/or configuration of a processor, and/or a computing system including a processor, that can execute programs of a compiler, such as compiler **1108**, and/or an analysis assistant, such as analysis assistant **1150**. Processor **1104A** is shown coupled to compiler **1108** by means of interface **1106A** and processor **1104B** is shown coupled to analysis assistant **1150** by means of interface **1106B**. Interface **1106A** can be, for example, an API for executing compiler **1108** on processor **1104A**, for compiler **1108** to receive input data, such as APP **1102**, and/or for compiler **1108** to output data, such as a graph and/or an alternative IR, of a dataflow program incorporating compiler high level decisions.

(196) Interface **1106B** can be, for example, an API for executing analysis assistant **1150** on processor **1104B**, for analysis assistant **1150** to receive input data, such as input data from compiler **1108** and/or developer **1160**, and/or for analysis assistant **1150** to output data, such as a inefficiency results and/or suggestions to developer **1160**, a storage device or system, compiler **1108**, and/or a computing system communicatively coupled to analysis assistant **1150** or a computing system on which analysis assistant **1150** can execute. Processors **1104A** and **1104B** can be the same processor, or can be different processors. Interfaces **1106A** and **1106B** can be the same interfaces and/or can be different interfaces and can comprise a variety of interfaces, such as APIs, CLIs, communications interfaces, data buses, I/O buses or links, network interfaces, and so forth.

(197) In implementations, compiler **1108** can be, compiler for compiling a dataflow program to execute on a dataflow computing system, such as a CGRS compiler for compiling a dataflow

program for execution on a CGRS, and/or on CGR hardware of a CGRS. Compiler **1108** can comprise, for example, a compiler that can perform methods similar or equivalent to method **500** of FIG. 5, method **900** of FIG. 9, and/or method **1000** of FIG. 10. Compiler **1108** can compile APP **1102** for execution on CGR hardware.

(198) Compiler **1108** is shown in FIG. 11 comprising, HL compiler **1120**, HLR **1130**, LL compiler **1112**, LLIR **1116**, HW SPEC **1114**, and analyzer **1140**. In implementations, HLR **1130** can comprise a graph corresponding to APP **1102** and/or LLIR **1116** can comprise an LLIR such as previously described in reference to FIGS. 6 and 7. HLR **1130** and LLIR **1116** can individually or, in combination, comprise an IR representing a set of operators of a dataflow program, such as a meta-pipe, and/or an IR for processing by an efficiency analyzer, such as in the example of method **900** in FIG. 9.

(199) Compiler **1108** can receive APP **1102** (e.g., via interface **1106A** or, another interface of compiler **1108** not shown in FIG. 11, such as an API for compiler **1108** to receive or to otherwise access APP **1102**) and compiler **1108** can generate HLR **1130** based on APP **1102**. Alternatively, compiler **1108** can receive HLR **1130**, such as from a storage device or memory storing HLR **1130** (e.g., storage or memory storing HLR **1130** from a prior creation of HLR **1130**). HLR **1130** can comprise an auxiliary graph, and/or can be generated by compiler **1108**. HW SPEC **1114** can comprise a hardware description of CGRS hardware, such as reconfigurable processors, and/or configurations of CGRS hardware.

(200) HL compiler **1120** can be a component (of, function) of compiler **1108** that can compile a dataflow program to an HLR, such as HLR **1130**. Such an HLR can include high level execution decisions mapping a dataflow program to CGR hardware for execution. HL compiler **1120** can comprise high level decision components (or, functions) to perform tiling, parallelizing, sectioning, and/or scheduling of operations of APP **1102**, which can individually and/or in combination determine high level execution alternatives and/or decisions. HL compiler **1120** can utilize a hardware description, such as HW SPEC **1114**, to assist in determining high level execution decisions.

(201) HLR **1130** is, correspondingly, shown in FIG. 11 comprising nodes **1132** and HL decisions **1134**. In implementations, Nodes **1132** can represent or correspond to operators, and/or operands/results of operators, of APP **1102**. HLR **1130** can describe a topology of nodes **1132** (e.g., a topology of operators, and/or dataflow among the operators, of APP **1102**). HL decisions **1134** can comprise high level execution alternatives and/or decisions, such as high level execution alternatives/decisions computed by HL compiler **1120**.

(202) LL compiler **1112** can comprise a component (or, function) of compiler **1108** that can compile an HLR to an LLIR. For example, in FIG. 11 LL compiler **1112** can compile HL compiler **1120** to LLIR **1116**. LL compiler **1112** can analyze LLIR **1116** to evaluate efficiency metrics of compiling HLR **1130** to low level hardware constructs. In combination, HL compiler **1120** and LL compiler **1112** can perform iterative evaluations of high level execution decisions and can determine more, and less, optimal high level decisions among them. In implementations, compiler **1108**, HL compiler **1120**, and LL compiler **1112** can be software elements (programs), hardware elements (e.g., special purpose compilation and/or graph processors or circuits), or a combination of software and hardware elements.

(203) In implementations analyzer **1140** can be an application analyzer such as previously described. For example, analyzer **1140** can perform a method similar or equivalent to method **900**, and/or operations thereof, and can analyze APP **1102** (e.g., as represented by HLR **1130** and/or LLIR **1116**) to predict hardware execution of APP **1102** on CGR hardware and/or determine application program inefficiencies. In FIG. 11, analyzer **1140** is shown comprising metrics **1142**, categories **1144**, and results **1146**. Metrics **1142** include, for example, program execution metrics and/or inefficiency metrics. Categories **1144** can include, for example, inefficiency categories and/or inefficiencies such as previously described, and results **1146** can comprise, for example,

inefficiency results produced by analyzing APP **1102** for inefficiencies associated with execution on particular CGR hardware, such as CGR hardware described in HW SPEC **1114**.

(204) In FIG. **11**, analysis assistant **1150** can be an analysis assistant such as previously described. For example, analysis assistant **1150** can perform a method similar or equivalent to method **1000** of FIG. **10**, and/or operations thereof. Analysis assistant **1150** can, for example, present results **1146** (or, alternative inefficiency results) determined by analyzer **1140**. Analysis assistant **1150** can present inefficiency results and/or suggestions to developer **1160**, and can interact (e.g., via an interface similar or equivalent to interface **1106B**) with developer **1160**.

(205) FIG. **11** shows analysis assistant **1150** comprising graphical data **1152**, tabular data **1154**, and optional suggestions **1156**. In implementations, optional suggestions **1156** can be suggestions and/or recommendations to improve inefficiencies determined by an analyzer, such as such as results **1146**, determined by an analyzer, such as analyzer **1140**. Graphical data **1152** can comprise graphical data (and/or, data to generate graphical output) corresponding to inefficiency results, such as metrics **1142** and/or results **1146**, determined by an analyzer, such as analyzer **1140**.

(206) Graphical data **1152** can comprise graphical data (and/or, data to generate graphical output) corresponding to inefficiency results, such as program execution metrics, inefficiency metrics, and/or inefficiencies. Tabular data **1154** can comprise tabular data (and/or, data to generate tabular output) corresponding to inefficiency results, such as program execution metrics, inefficiency metrics, and/or inefficiencies, such as can be included in metrics **1142** and/or results **1146**. Tabular data **1154** can comprise tabular data (and/or, data to generate tabular output) corresponding to suggestions and/or recommendations, such as optional suggestions **1156**.

(207) While FIG. **11** depicts analyzer **1140** as included in compiler **1108**, and analysis assistant **1150** as a component of a system separate from compiler **1108**, this is for purposes only of illustrating the disclosure and not intended to limit implementations. It would be apparent to one of ordinary skill in the art that, in implementations, analyzer **1140** need not be a component of compiler **1108**, that analysis assistant **1150** can be a component of compiler **1108**, and/or that any or all of compiler **1108**, analyzer **1140**, and analysis assistant **1150** can be implemented in any combination and/or structural arrangement of computing systems and/or components thereof. It would be equally apparent to one of ordinary skill in the art that such a system can comprise a variety of components, and/or arrangements of components, alternative to the example of FIG. **11**, that such alternative components/arrangements can perform iterative compilation methods and/or operations of the disclosure, and that such alternative implementations will be within the scope and spirit of the disclosure.

(208) While not shown explicitly in FIG. **11**, in computing system **1100** HW SPEC **1114**, HLR **1130**, LLIR **1116**, metrics **1142**, categories **1144**, results **1146**, graphical data **1152**, tabular data **1154**, and/or optional suggestions **1156**, and/or components of these, can be stored in a storage medium and/or memory included (but, not shown explicitly) in, or coupled to, computing system **1100**, and/or accessible to processors **1104A** and/or **1104B**. Programming instructions for compiler **1108**, HL compiler **1120**, LL compiler **1112**, analyzer **1140**, and/or analysis assistant **1150**, such as to perform methods and/or operations of the disclosure, can be stored in a storage medium and/or memory included (but, not shown explicitly) in, or coupled to, computing system **1100**, and/or accessible to processors **1104A** and/or **1104B**.

(209) In implementations, a storage medium can comprise any form of persistent memory or medium, such as a disk or optical storage device, or a storage medium of a computer storage subsystem. A memory can comprise any computer memory that can be accessible to compiler **1108**, and/or processors **1104A** and/or **1104B** while executing programs of compiler **1108**, analyzer **1140**, and/or analysis assistant **1150**. For example, while not shown explicitly in FIG. **11**, computing system **1100** can include a main memory, on-chip and/or off-chip memories of a CGRA, cache memories, and/or memories of a CGRS (e.g., PMUs, DRAMs, and/or SRAMs).

(210) In implementations, an apparatus including a memory and at least one processor coupled to

the memory, can perform example methods, and/or operations thereof, of the disclosure. Implementations can comprise a means for carrying out example methods, and/or operations thereof, of the disclosure; the means can include (i) hardware module(s), (ii) software module(s) executing on one or more hardware processors, or (iii) a combination of hardware and software modules; any of (i)-(iii) configured to implement example methods, and/or operations thereof, of the disclosure, and in which the software modules are stored in a computer readable storage medium (or multiple such media).

(211) Implementations can comprise a computer program product and can include a computer readable storage medium (or media) having computer readable program instructions of the computer program product incorporated therein. It will be understood by one of ordinary skill in the art that computer readable program instructions can implement each or any combination of operations and/or structure of the disclosure, such as illustrated by the drawings and described herein.

(212) The computer readable program instructions can be provided to one or more processors, and/or other elements, of a computing system or apparatus to produce a machine which can execute, via the processor(s), to implement operations and/or actions similar or equivalent to those of the disclosure. The computer readable program instructions can be stored in a computer readable storage medium that can direct one or more processors, and/or other elements, of a computing system or apparatus to function in a particular manner, such that the computer readable storage medium comprises an article of manufacture including instructions to implement operations and/or structures similar or equivalent to those of the disclosure.

(213) The computer readable program instructions of the computer program product can cause one or more processors to perform operations of the disclosure. A sequence of program instructions, and/or an assembly of one or more interrelated programming modules, of the computer program product can direct one or more one or more processors and/or computing elements of a computing system to implement the elements and/or operations of the disclosure including, but not limited to, the structures and operations illustrated and/or described in the present disclosure.

(214) A computer readable storage medium can comprise any tangible (e.g., hardware) device, or combination of tangible devices, that can store instructions of the computer program product and that can be read by a computing element to download the instructions for use by a processor. A computer readable storage medium can comprise, but is not limited to, electronic, magnetic, optical, electromagnetic, and/or semiconductor storage devices, or any combination of these. A computer readable storage medium can comprise a portable storage medium, such as a magnetic disk/diskette, optical disk (CD or DVD); a volatile and/or non-volatile memory; a memory stick, a mechanically encoded device, and any combination of these. A computer readable storage medium, as used herein, is not to be construed as being transitory signals per se, such as electrical signals transmitted through a wire, radio waves or other freely propagating electromagnetic waves, or electromagnetic waves propagating through a wave transmission medium (e.g., a wave guide or fiber-optic cable).

(215) The computer readable program instructions can be communicated from the computer readable storage medium to the one or more computing/processing devices, via a programming API of a computing system, and/or a communications interface of a computing system, having access to the computer readable storage medium, and/or a programming API of a computing system, and/or a communications interface of the one or more computing/processing devices. The API(s) and/or communications interface(s) can couple communicatively and/or operatively to a network, such as the Internet, a local area network, a wide area network, and/or a wireless network. The API(s) and/or communications interface(s) can receive the computer readable program instructions read from computer readable storage medium and can forward the computer readable program instructions to the one or more computing/processing devices via the API(s), communications interface(s), and/or network.

(216) The computer readable program instructions of the computer program product can comprise machine language and/or assembly language instructions, instruction-set-architecture (ISA) instructions, microcode and/or firmware instructions, state-setting data, configuration data for integrated circuitry, source code, and/or object code. The instructions and/or data can be written in any combination of one or more programming languages.

(217) The computer readable program instructions can execute entirely, or in part, on a user's computer, as a stand-alone software package; partly on a user's computer and partly on a remote computer; or, entirely on a remote computer. A remote computer can be connected to a user's computer through any type of network, including a local area network (LAN) or a wide area network (WAN). In implementations, electronic circuitry including, for example, FPGA, PLAs, and or CGRPs can execute the computer readable program instructions by utilizing state information of the computer readable program instructions to configure the electronic circuitry to perform operations or elements of the disclosure, such as illustrated by the drawings and described herein.

(218) In implementations, computer readable program instructions can also be loaded onto a computing system, or component(s) thereof, to cause the computing system and/or component(s) thereof to perform a series of operational operations to produce a computer implemented process, such that the instructions which execute on the computing system, or component(s) thereof, implement the operations or elements of the disclosure, such as illustrated by the drawings and described herein.

(219) The flowchart and block diagrams in the Drawings and Incorporations illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various implementations of the present invention. Individual elements illustrated in the Figures—such as individual operations illustrated in the flowcharts or individual blocks of block diagrams—can represent a module, segment, or portion of executable instructions for implementing the disclosed function(s). In various alternative implementations, particular operations can occur in an order differing from that illustrated in the examples of the drawings. For example, two operations shown in succession in a diagram of the disclosure may, in a particular implementation, be executed substantially concurrently, or can sometimes be executed in a reverse order, depending upon the functionality involved. It will be further noted that particular blocks of the block diagrams, operations of the flowchart illustrations, and/or combinations of blocks in the block diagrams and/or flowchart illustrations, can be implemented using special purpose hardware and/or systems that, individually or in combination, perform the specified functions, acts, and/or computer instructions.

(220) In implementations, features of the disclosure, and/or elements of features of the disclosure, can be combined. In the interest of conciseness, such combinations are not individually enumerated and are not repeated with each base set of features. One of ordinary skill in the art will appreciate how various individual features of the disclosure can be readily combined with sets of base features identified in various example implementations. Features of the disclosure, and/or elements of features of the disclosure, are not meant to be mutually exclusive, exhaustive, or restrictive in implementations. Further, the disclosure is not limited to these features but rather encompasses all possible combinations, modifications, equivalents, and variations within the scope of the disclosure.

(221) Terminology used herein, and the examples disclosed, are chosen to illustrate the principles of the implementations, the practical application or technical improvement over alternative technologies, and to enable others of ordinary skill in the art to understand the implementations disclosed herein. The disclosure illustrates various example implementations, and the examples are intended to illustrate principles and aspects of the disclosure, but are not intended to limit implementations, nor intended to be exhaustive of implementations that can be conceived within the scope of the disclosure. It would be apparent to one of ordinary skill in the art that alternative implementations can comprise modifications and combinations within the spirit of the disclosure

and the scope of the claims.

(222) As can be seen in the foregoing examples, features of the disclosure can comprise methods and apparatus of computing systems. A summary of example implementations of such features includes:

Example Implementation 1

(223) A method comprises: selecting, by a computer-implemented efficiency analyzer, from an intermediate representation of a dataflow program, a set of operators of the dataflow program, the set of operators included in a first mapping of the set of operators to first hardware to execute the set of operators, the first mapping included in the intermediate representation, the first hardware included in a first computing system for executing the dataflow program; computing, by the efficiency analyzer, based on a hardware description and the first mapping, a first predicted execution metric associated with executing the set of operators on the first hardware, the hardware description corresponding to hardware included in the first computing system; determining, by the efficiency analyzer, based on the first predicted execution metric, a first inefficiency metric associated with the executing the set of operators on the first hardware; determining, by the efficiency analyzer, based on the hardware description and the first inefficiency metric, a first inefficiency associated with the dataflow program; and, outputting, by the efficiency analyzer, first inefficiency results comprising the first inefficiency.

Example Implementation 2

(224) The example of implementation 1, wherein the method is included in a method for a compiler to compile the dataflow program to execute on the first computing system.

Example Implementation 3

(225) The example of implementation 1, wherein the first computing system comprises a coarse-grain reconfigurable computing system.

Example Implementation 4

(226) The example of implementation 1, the method further comprising determining, by the compiler, based on the evaluating the first trial result, a modification of the first high level decision.

Example Implementation 5

(227) The example of implementation 1, wherein the method further comprises: determining, by the inefficiency analyzer, that the set of operators comprises a nested meta-pipe; selecting, by the efficiency analyzer, from the intermediate representation, a second mapping comprising a mapping of the nested meta-pipe to second hardware to execute operators of the nested meta-pipe, the second hardware included in the first computing system; computing, by the efficiency analyzer, based on the hardware description and the second mapping, a second predicted execution metric associated with executing the operators of the nested meta-pipe on the second hardware; determining, by the efficiency analyzer, based on the second predicted execution metric, a second inefficiency metric associated with the executing the operators of the nested meta-pipe on the second hardware; determining, by the efficiency analyzer, based on the hardware description and the second inefficiency metric, a second inefficiency associated with the dataflow program; and, outputting, by the efficiency analyzer, second inefficiency results comprising the second inefficiency.

Example Implementation 6

(228) The example of implementation 1, wherein the first predicted execution metric comprises a predicted execution latency to execute an operator among the set of operators on the first hardware; and, wherein the efficiency analyzer determining the first inefficiency metric comprises: computing, by the efficiency analyzer, a theoretical minimum execution latency to execute the operator; and, determining, by the efficiency analyzer, the first inefficiency metric based on a difference between the predicted execution latency and the theoretical minimum execution latency.

Example Implementation 7

(229) The example of implementation 1, wherein the method of the outputting, by the

efficiency analyzer, the first inefficiency results further comprises outputting, by the first efficiency analyzer, the first inefficiency results to a component of a second computing system, the component selected from the group consisting of an application programming interface of a compiler; an application programming interface of the second computing system; a memory of the second computing system; a storage device of the second computing system; and, an analysis assistant program of the second computing system.

Example Implementation 8

(230) The example of implementation 1, wherein the first predicted execution metric is selected from the group consisting of a number of processors of the first computing system to execute the first mapping; a number of parallel computations to execute the first mapping; a latency of a first processor of the first computing system to execute the first mapping; a utilization of a second processor of the first computing system to execute the first mapping; a number of stage buffers of the first computing system required to execute the first mapping; an amount of a first memory of the first computing system required to execute the first mapping; a size of a second memory of the first computing system to execute the first mapping; a memory transfer latency of the first computing system to execute the first mapping; a utilization of a third memory of the first computing system to execute the first mapping; a utilization of a particular type of memory of the first computing system to execute the first mapping; and, a utilization of a particular memory of the first computing system to execute the first mapping.

Example Implementation 9

(231) The example of implementation 1, wherein an inefficiency metric included in the first inefficiency metric is selected from the group consisting of initialization of an operator, included in the set of operators, to execute on the first computing system; initialization of a first hardware pipeline of the first computing system to execute the operator; termination of a second hardware pipeline of the first computing system to execute the operator; completing remaining computations of the operator by a processor of the first computing system; an inefficiency in a low level program for executing the operator on the first computing system; an inefficiency in a sequence of low level programs for executing the first mapping on the first computing system; and, a difference between an actual execution latency for executing the first mapping on the first computing system and a theoretical minimum execution latency for executing the first mapping on the first computing system.

Example Implementation 10

(232) A computer program product comprises a computer readable storage medium having program instructions embodied therewith, wherein the program instructions are executable by at least one processor of a first computing system to cause the at least one processor to: select, from an intermediate representation of a dataflow program, a set of operators of the dataflow program, the set of operators included in a first mapping of the set of operators to first hardware to execute the set of operators, the first mapping included in the intermediate representation, the first hardware included in a first computing system for executing the dataflow program; compute, based on a hardware description and the first mapping, first predicted execution metric associated with executing the set of operators on the first hardware, the hardware description corresponding to hardware included in the first computing system; determine, based on the first predicted execution metric, first inefficiency metric associated with the executing the set of operators on the first hardware; determine, based on the hardware description and the first inefficiency metric, a first inefficiency associated with the dataflow program; and, output first inefficiency results comprising the first inefficiency.

Example Implementation 11

(233) The example of implementation 10, wherein the program instructions are executable by the at least one processor to further cause the at least one processor to: determine that the set of operators comprises a nested meta-pipe; select, from the intermediate representation, a second mapping

comprising a mapping of the nested meta-pipe to second hardware to execute operators of the nested meta-pipe, the second hardware included in the first computing system; compute, based on the hardware description and the second mapping, a second predicted execution metric associated with executing the operators of the nested meta-pipe on the second hardware; determine, based on the second predicted execution metric, a second inefficiency metric associated with the executing the operators of the nested meta-pipe on the second hardware; determine, based on the hardware description and the second inefficiency metric, a second inefficiency associated with the dataflow program; and, output second inefficiency results comprising the second inefficiency.

Example Implementation 12

(234) A first computing system comprises: at least one processor included in the first computing system; an intermediate representation of a dataflow program; a hardware description corresponding to hardware included in a second computing system, the second computing system configurable to execute the dataflow program; and, an efficiency analyzer.

(235) The efficiency analyzer is configured to execute on the at least one processor to: select, from the intermediate representation, a set of operators of the dataflow program, the set of operators included in a first mapping of the set of operators to first hardware to execute the set of operators, the first mapping included in the intermediate representation, the first hardware among the hardware included in the second computing system; compute, based on the hardware description and the first mapping, a first predicted execution metric associated with executing the set of operators on the first hardware; determine, based on the first predicted execution metric, a first inefficiency metric associated with the executing the set of operators on the first hardware; determine, based on the hardware description and the first inefficiency metric, a first inefficiency associated with the dataflow program; and, output first inefficiency results comprising the first inefficiency.

Example Implementation 13

(236) The example of implementation 12, wherein the efficiency analyzer is further configured to execute on the at least one processor to: determine that the set of operators comprises a nested meta-pipe; select, from the intermediate representation, a second mapping comprising a mapping of the nested meta-pipe to second hardware to execute operators of the nested meta-pipe, the second hardware included in the first computing system; compute, based on the hardware description and the second mapping, a second predicted execution metric associated with executing the operators of the nested meta-pipe on the second hardware; determine, based on the second predicted execution metric, a second inefficiency metric associated with the executing the operators of the nested meta-pipe on the second hardware; determine, based on the hardware description and the second inefficiency metric, a second inefficiency associated with the dataflow program; and, output second inefficiency results comprising the second inefficiency.

Example Implementation 14

(237) The example of implementation 12, wherein the first predicted execution metric comprises a predicted execution latency to execute an operator among the set of operators on the first hardware; and, wherein the efficiency analyzer configured to execute on the at least one processor to determine the first inefficiency metric comprises the efficiency analyzer further configured to execute on the at least one processor to: compute a theoretical minimum execution latency to execute the operator; and, determine the first inefficiency metric based on a difference between the predicted execution latency and the theoretical minimum execution latency.

Example Implementation 15

(238) The example of implementation 12, wherein the efficiency analyzer is further configured to execute on the at least one processor to: determine an inefficiency category associated with the first inefficiency; and, output, among the first inefficiency results, the inefficiency category.

Example Implementation 16

(239) The example of implementation 12, wherein the efficiency analyzer configured to execute on

the at least one processor to output the first inefficiency results comprises the efficiency analyzer further configured to execute on the at least one processor to output the first inefficiency results to a component of a third computing system, the component selected from the group consisting of an application programming interface of a compiler; an application programming interface of the third computing system; a memory of the third computing system; a storage device of the third computing system; and, an analysis assistant program of the third computing system.

Example Implementation 17

(240) The example of implementation 12, wherein the first predicted execution metric is selected from the group consisting of a number of processors of the second computing system to execute the first mapping; a number of parallel computations to execute the first mapping; a latency of a first processor of the second computing system to execute the first mapping; a utilization of a second processor of the second computing system to execute the first mapping; a number of stage buffers of the second computing system required to execute the first mapping; an amount of a first memory of the second computing system required to execute the first mapping; a size of a second memory of the second computing system to execute the first mapping; a memory transfer latency of the second computing system to execute the first mapping; a utilization of a third memory of the second computing system to execute the first mapping; a utilization of a particular type of memory of the second computing system to execute the first mapping; and, a utilization of a particular memory of the second computing system to execute the first mapping.

Example Implementation 18

(241) The example of implementation 12, wherein the first inefficiency metric is selected from the group consisting of initialization of an operator, included in the set of operators, to execute on the second computing system; initialization of a first hardware pipeline of the second computing system to execute the operator; termination of a second hardware pipeline of the second computing system to execute the operator; completing remaining computations of the operator by a processor of the second computing system; an inefficiency in a low level program for executing the operator on the second computing system; an inefficiency in a sequence of low level programs for executing the first mapping on the second computing system; and, a difference between an actual execution latency for executing the first mapping on the second computing system and a theoretical minimum execution latency for executing the first mapping on the second computing system.

Example Implementation 19

(242) The example of implementation 12, wherein the efficiency analyzer is associated with a compiler for compiling the dataflow program to execute on the second computing system.

Example Implementation 20

(243) The example of implementation 12, wherein the second computing system comprises a coarse-grain reconfigurable computing system.

Example Implementation 21

(244) A method comprises: initiating, by a computer-implemented analysis assistant, responsive to an interface of a first computing system, a presentation session comprising a presentation of inefficiency results associated with a mapping of a set of operators of a dataflow program to execute the set of operators on hardware of a second computing system, the analysis assistant included in the first computing system; receiving, by the analysis assistant, an inefficiency included among the inefficiency results; composing by the analysis assistant, formatted inefficiency results comprising a presentation format of the inefficiency to assist a developer of the dataflow program to interpret the inefficiency; and, outputting, by the analysis assistant, to an interface of a third computing system, the formatted inefficiency results, the interface of the third computing system comprising an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.

Example Implementation 22

(245) The example of implementation 21, wherein the method further comprises: determining, by

the analysis assistant, a suggestion to modify the dataflow program to improve the inefficiency; and, outputting, by the analysis assistant, using a second output interface of the third computing system, the suggestion.

Example Implementation 23

(246) The example of implementation 22, wherein the suggestion to modify the dataflow program to improve the inefficiency comprises one of a suggestion to modify the dataflow program to reduce the inefficiency and a suggestion to modify the dataflow program to remove the inefficiency.

Example Implementation 24

(247) The example of implementation 21, wherein the presentation session comprises an interactive presentation of the formatted inefficiency results.

Example Implementation 25

(248) The example of implementation 24, wherein the interactive presentation is included in an interaction of a compiler with a developer to program the dataflow program.

Example Implementation 26

(249) The example of implementation 21, wherein the second computing system comprises a coarse-grain reconfigurable system.

Example Implementation 27

(250) The example of implementation 21, wherein the hardware of the second computing system comprises at least one reconfigurable processor;

(251) wherein the mapping comprises executing an operator, among the set of operators, on a processor among the at least one reconfigurable processor; and wherein the inefficiency results are further associated with the executing the operator on the processor among the at least one reconfigurable processor.

Example Implementation 28

(252) The example of implementation 21, wherein at least one of the interface of the first computing system and the interface of the third computing system comprises an interface of a compiler for compiling the dataflow program.

Example Implementation 29

(253) The example of implementation 21, wherein the presentation format is selected from the group consisting of a tabular format; a graphical format; a pie chart; a bar chart; a latency waterfall chart; and, a combination of a bar chart and a latency waterfall chart.

Example Implementation 30

(254) The example of implementation 21, wherein the inefficiency is included among an inefficiency category selected from the group consisting of an inefficiency in an operator program to execute an operator of the dataflow program; a pipeline imbalance; a memory stall; a transient effect; an unused hardware component of the second computing system; and, an underutilized hardware component of the second computing system.

Example Implementation 31

(255) A computer program product comprises computer readable storage medium having program instructions embodied therewith, wherein the program instructions are executable by at least one processor to cause the at least one processor to: initiate, responsive to an interface of a first computing system, a presentation session comprising a presentation of inefficiency results associated with a mapping of a set of operators of a dataflow program to execute the set of operators on hardware of a second computing system; receive an inefficiency included among the inefficiency results; compose formatted inefficiency results comprising a presentation format of the inefficiency to assist a developer of the dataflow program to interpret the inefficiency; and, output, to an interface of a third computing system, the formatted inefficiency results, the interface of the third computing system comprising an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.

Example Implementation 32

(256) The example of implementation 31, wherein the program instructions are executable by the at least one processor to further cause the at least one processor to: determine, based on the inefficiency, a suggestion to modify the dataflow program to improve the inefficiency; and, output, using a second output interface of the third computing system, the suggestion.

Example Implementation 33

(257) A first computing system comprises: a first processor; a second computing system comprising a second processor; an efficiency analyzer configured to execute on the second processor to determine inefficiency results associated with executing a dataflow program on a third computing system; and, an analysis assistant.

(258) The analysis assistant is configured to execute on the first processor to: initiate, responsive to an interface of the first computing system, a presentation session comprising a presentation of the inefficiency results; receive the inefficiency results; select, from among the inefficiency results, an inefficiency; compose formatted inefficiency results comprising a presentation format of the inefficiency to assist a developer of the dataflow program to interpret the inefficiency; and, output, to an interface of a fourth computing system, the formatted inefficiency results, the interface of the fourth computing system comprising an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.

Example Implementation 34

(259) The example of implementation 33, wherein the analysis assistant is configured to execute on the first processor to further: determine a suggestion to modify the dataflow program to improve the inefficiency; and, output, using a second output interface of the fourth computing system, the suggestion.

Example Implementation 35

(260) The example of implementation 33, wherein the analysis assistant is configured to execute on the first processor to initiate, responsive to an interface of the first computing system, the presentation session comprising the analysis assistant configured to execute on the first processor to further initiate the presentation session to comprise an interactive presentation of the formatted inefficiency results.

Example Implementation 36

(261) The example of implementation 35, wherein the interactive presentation is included in an interaction of a compiler with a developer to program the dataflow program.

Example Implementation 37

(262) The example of implementation 33, wherein the presentation format is selected from the group consisting of a tabular format; a graphical format; a pie chart; a bar chart; a latency waterfall chart; and, a combination of a bar chart and a latency waterfall chart.

Example Implementation 38

(263) The example of implementation 33, wherein the third computing system comprises at least one reconfigurable processor; and, wherein the inefficiency results associated with the executing the dataflow program on the third computing system comprise inefficiency results associated with the executing an operator of the dataflow program on a processor among the at least one reconfigurable processor.

Example Implementation 39

(264) The example of implementation 33, wherein at least one of the interface of the first computing system and the interface of the third computing system comprises an interface of a compiler for compiling the dataflow program.

Example Implementation 40

(265) The example of implementation 33, wherein the inefficiency is included among an inefficiency category selected from the group consisting of an inefficiency in an operator program to execute an operator of the dataflow program; a pipeline imbalance; a memory stall; a transient

effect; an unused hardware component of the second computing system; and, an underutilized hardware component of the second computing system.

Claims

1. A method, the method comprising: initiating, by a computer-implemented analysis assistant, responsive to an interface of a first computing system, a presentation session comprising a presentation of inefficiency results associated with a mapping of a set of operators of a dataflow program to execute the set of operators on hardware of a second computing system, the analysis assistant included in the first computing system; receiving, by the analysis assistant, an inefficiency included among the inefficiency results; composing by the analysis assistant, formatted inefficiency results comprising the inefficiency having a presentation format to assist a developer of the dataflow program to interpret the inefficiency; and, outputting, by the analysis assistant, to an interface of a third computing system, the formatted inefficiency results, the interface of the third computing system comprising at least a portion to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.
2. The method of claim 1, wherein the method further comprises: determining, by the analysis assistant, a suggestion to modify the dataflow program to improve the inefficiency; and, outputting, by the analysis assistant, using a second output interface of the third computing system, the suggestion.
3. The method of claim 2, wherein the suggestion to modify the dataflow program to improve the inefficiency comprises one of a suggestion to modify the dataflow program to reduce the inefficiency and a suggestion to modify the dataflow program to remove the inefficiency.
4. The method of claim 1, wherein the presentation session comprises an interactive presentation of the formatted inefficiency results.
5. The method of claim 4, wherein the interactive presentation is included in an interaction of a compiler with a developer to program the dataflow program.
6. The method of claim 1, wherein the second computing system comprises a coarse-grain reconfigurable system.
7. The method of claim 1, wherein the hardware of the second computing system comprises at least one reconfigurable processor; wherein the mapping comprises executing an operator, among the set of operators, on a processor among the at least one reconfigurable processor; and wherein the inefficiency results are further associated with the executing the operator on the processor among the at least one reconfigurable processor.
8. The method of claim 1, wherein at least one of the interface of the first computing system and the interface of the third computing system comprises an interface of a compiler for compiling the dataflow program.
9. The method of claim 1, wherein the presentation format is selected from the group consisting of a tabular format; a graphical format; a pie chart; a bar chart; a latency waterfall chart; and, a combination of a bar chart and a latency waterfall chart.
10. The method of claim 1, wherein the inefficiency is included among an inefficiency category selected from the group consisting of an inefficiency in an operator program to execute an operator of the dataflow program; a pipeline imbalance; a memory stall; a transient effect; an unused hardware component of the second computing system; and, an underutilized hardware component of the second computing system.
11. A computer program product, the computer program product comprising a computer readable storage medium having program instructions embodied therewith, wherein the program instructions are executable by at least one processor to cause the at least one processor to: initiate, responsive to an interface of a first computing system, a presentation session comprising a presentation of inefficiency results associated with a mapping of a set of operators of a dataflow program to

execute the set of operators on hardware of a second computing system; receive an inefficiency included among the inefficiency results; compose formatted inefficiency results comprising the inefficiency having a presentation format to assist a developer of the dataflow program to interpret the inefficiency; and, output, to an interface of a third computing system, the formatted inefficiency results, the interface of the third computing system comprising at least a portion an interface to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.

12. The computer program product of claim 11, wherein the program instructions are executable by the at least one processor to further cause the at least one processor to: determine, based on the inefficiency, a suggestion to modify the dataflow program to improve the inefficiency; and, output, using a second output interface of the third computing system, the suggestion.

13. A first computing system, the first computing system comprising: a first processor; a second computing system comprising a second processor; an efficiency analyzer configured to execute on the second processor to determine inefficiency results associated with executing a dataflow program on a third computing system; and, an analysis assistant configured to execute on the first processor to: initiate, responsive to an interface of the first computing system, a presentation session comprising a presentation of the inefficiency results; receive the inefficiency results; select, from among the inefficiency results, an inefficiency; compose formatted inefficiency results comprising the inefficiency having a presentation format to assist a developer of the dataflow program to interpret the inefficiency; and, output, to an interface of a fourth computing system, the formatted inefficiency results, the interface of the fourth computing system comprising at least a portion to output the formatted inefficiency results for use by the developer to improve the dataflow program in association with the inefficiency.

14. The first computing system of claim 13, wherein the analysis assistant is configured to execute on the first processor to further: determine a suggestion to modify the dataflow program to improve the inefficiency; and, output, using a second output interface of the fourth computing system, the suggestion.

15. The first computing system of claim 13, wherein the analysis assistant configured to execute on the first processor to initiate, responsive to an interface of the first computing system, the presentation session comprises the analysis assistant configured to execute on the first processor to further initiate the presentation session to comprise an interactive presentation of the formatted inefficiency results.

16. The first computing system of claim 15, wherein the interactive presentation is included in an interaction of a compiler with a developer to program the dataflow program.

17. The first computing system of claim 13, wherein the presentation format is selected from the group consisting of a tabular format; a graphical format; a pie chart; a bar chart; a latency waterfall chart; and, a combination of a bar chart and a latency waterfall chart.

18. The first computing system of claim 13, wherein the third computing system comprises at least one reconfigurable processor; and, wherein the inefficiency results associated with the executing the dataflow program on the third computing system comprise inefficiency results associated with the executing an operator of the dataflow program on a processor among the at least one reconfigurable processor.

19. The first computing system of claim 13, wherein at least one of the interface of the first computing system and the interface of the third computing system comprises an interface of a compiler for compiling the dataflow program.

20. The first computing system of claim 13, wherein the inefficiency is included among an inefficiency category selected from the group consisting of an inefficiency in an operator program to execute an operator of the dataflow program; a pipeline imbalance; a memory stall; a transient effect; an unused hardware component of the second computing system; and, an underutilized hardware component of the second computing system.

