

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12411672
B2
September 09, 2025
Koupy; Petr et al.

Method of distributed graph loading for minimal communication and good balance via lazy materialization and directory indirection using indexed tabular representation

Abstract

Techniques herein minimally communicate between computers to repartition a graph. In embodiments, each computer receives a partition of edges and vertices of the graph. For each of its edges or vertices, each computer stores an intermediate representation into an edge table (ET) or vertex table. Different edges of a vertex may be loaded by different computers, which may cause a conflict. Each computer announces that a vertex resides on the computer to a respective tracking computer. Each tracking computer makes assignments of vertices to computers and publicizes those assignments. Each computer that loaded conflicted vertices transfers those vertices to computers of the respective assignments. Each computer stores a materialized representation of a partition based on: the ET and vertex table of the computer, and the vertices and edges that were transferred to the computer. Edges stored in the materialized representation are stored differently than edges stored in the ET.

Inventors: Koupy; Petr (Blansko, CZ), Manhardt; Thomas (Foster City, CA), Depner; Siegfried (Sunnyvale, CA), Hong; Sungpack (Palo Alto, CA), Chafi; Hassan (San Mateo, CA)

Applicant: Oracle International Corporation (Redwood Shores, CA)

Family ID: 65897248

Assignee: Oracle International Corporation (Redwood Shores, CA)

Appl. No.: 18/084406

Filed: December 19, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230121198 A1	Apr. 20, 2023

Related U.S. Application Data

division parent-doc US 17069104 20201013 US 11561780 child-doc US 18084406
division parent-doc US 15721848 20170930 US 10846069 20201124 child-doc US 17069104

Publication Classification

Int. Cl.: **G06F16/00** (20190101); **G06F8/41** (20180101); **G06F9/50** (20060101); **G06F11/34** (20060101); **G06F16/27** (20190101); **G06F17/10** (20060101); **G06F17/16** (20060101); G06F3/06 (20060101); G06F11/10 (20060101); G06F16/901 (20190101)

U.S. Cl.:

CPC **G06F8/453** (20130101); **G06F9/5038** (20130101); **G06F11/3433** (20130101); **G06F16/278** (20190101); **G06F17/10** (20130101); **G06F17/16** (20130101); G06F3/061 (20130101); G06F11/1076 (20130101); G06F16/9024 (20190101)

Field of Classification Search

CPC: G06F (16/278)

USPC: 707/798

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8321492	12/2011	Sengupta	708/620	G06F 17/16
2017/0068746	12/2016	Levin	N/A	N/A
2018/0039710	12/2017	Chen	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
WO2015099562	12/2014	WO	N/A

OTHER PUBLICATIONS

Presta et al., "Large-scale Graph Partitioning with Apache Giraph", Engineering Blog, Facebook Code, dated Apr. 9, 2014, 8 pages. cited by applicant

Karypis, George, Metis, A Software Package for Partitioning Unstructured Graphs, Partitioning Meshes, and Computing Fill-Reducing Orderings of Sparse Matrices Version 5.1.0, dated Mar. 30, 2013, 35pgs. cited by applicant

Gonzalez et al., "GraphX Processing in a Distributed Dataflow Framework", <https://amplab.cs.berkeley.edu/wp-content/uploads/2014/09/graphx.pdf>, dated 2014, 15 pages. cited by applicant

Primary Examiner: Daye; Chelcie L

Attorney, Agent or Firm: Hickman Becker Bingham Ledesma LLP

Background/Summary

BENEFIT CLAIM (1) This application claims the benefit, under 35 U.S.C. § 120, as a divisional of U.S. patent application Ser. No. 17/069,104, filed Oct. 13, 2020; which is a divisional of U.S. patent application Ser. No. 15/721,848, filed Sep. 30, 2017, now U.S. Pat. No. 10,846,069 issued Nov. 20, 2020, and the entire contents of these applications are hereby incorporated by reference as if fully set forth herein. The applicant(s) hereby rescind any disclaimer of claim scope in the parent application or the prosecution history thereof and advise the USPTO that the claims in this application may be broader than any claim in the parent application.

FIELD OF THE DISCLOSURE

(1) This disclosure relates to graph data partitioning for a distributed system. Presented herein are coordination techniques that need less communication and processing to achieve well balanced repartitioning.

BACKGROUND

(2) Graph processing is an analytic technique of increasing importance for Big Data.

Accommodating a huge dataset may involve horizontal scaling. Thus, distributed graph processing may be required and involve multiple computers connected by network.

(3) Although there are various distributed graph processing tools available, graph loading remains problematic. It takes too long to load graph data into the distributed graph processing system itself, before applying any analysis at all.

(4) The overhead of distributed graph loading typically arises from the process and consequences of assigning and retrieving data partitions. Partitioned loading entails multiple computers simultaneously retrieving data in non-overlapping and/or contentious ways. Data needs to be distributed evenly amongst computers. For instance, each computer should have approximately a same number of edges.

(5) Depending on the architecture of distributed graph processing system, there can be additional constraints. For example, directed edges originating from a same vertex should be loaded into a same computer.

(6) Distributed loading of a large graph typically requires multiple phases with complex computation and much network communication. Conventional distributed loading solutions entail assumptions and heuristics that reduce efficiency. For example, eager materialization of the initially loaded graph increases the computational cost of balancing (repartitioning). A straightforward way to partition a graph is to compute hash values of vertex keys and allocate vertices to computers based on those hash values. However, this approach requires much communication during loading and may still result in data imbalance when vertex degree (edge count) distribution is skewed by the graph content.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) In the drawings:

(2) FIG. 1 is a block diagram that depicts an example distributed system that uses minimal communication to repartition a distributed graph to achieve a balanced workload, in an embodiment;

(3) FIG. 2 is a scenario diagram that depicts an example process that uses minimal communication to repartition a distributed graph to achieve a balanced workload, in an embodiment;

(4) FIG. 3 is a block diagram that depicts an example computer that uses a segmented array to implement each column of an edge table or vertex table, in an embodiment;

- (5) FIG. 4 is a block diagram that depicts an example computer that optimally merges two segmented arrays, in an embodiment;
- (6) FIG. 5 is a block diagram that depicts an example computer that has various optimizations that accelerate the use of a vertex table or edge table, in an embodiment;
- (7) FIG. 6 is a block diagram that depicts an example computer cluster arranged in a logical topology that is dynamically managed by tracking computers, in an embodiment;
- (8) FIG. 7 is a scenario diagram that depicts example process that uses batching and iteration for deconfliction of redundantly loaded vertices, in an embodiment;
- (9) FIG. 8 is a scenario that depicts an example process that uses ghost vertices to avoid conflicts, in an embodiment;
- (10) FIG. 9 is a scenario that depicts an example process that uses the master computer to rebalance partitions, in an embodiment;
- (11) FIG. 10 is a block diagram that illustrates a computer system upon which an embodiment of the invention may be implemented;
- (12) FIG. 11 is a block diagram that illustrates a basic software system 500 that may be employed for controlling the operation of a computing system.

DETAILED DESCRIPTION

(13) In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present invention. It will be apparent, however, that the present invention may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form in order to avoid unnecessarily obscuring the present invention.

(14) Embodiments are described herein according to the following outline: 1.0 General Overview 2.0 Example Distributed System 2.1 Physical Topology 2.2 Initial Partition 2.3 Graph Retrieval 2.4 Vertex Table and Edge Table 2.5 Tracking Computer 2.6 Vertex Reassignment 2.7 Virtual Movement 2.8 Actual Movement 3.0 Example Repartitioning Process 4.0 Intermediate Representation 4.1 Segmented Array 5.0 Segments 5.1 Random Access 5.2 Optimal Merge 6.0 Table Optimizations 6.1 Indexing 6.2 Multithreading 6.3 Table View 7.0 Logical Topology 8.0 Deconfliction 9.0 Ghost Vertex 10.0 Rebalancing 11.0 Hardware Overview 12.0 Software Overview 13.0 Cloud Computing

1.0 General Overview

(15) Techniques are provided for minimal communication between a set of computers to repartition a distributed graph to achieve a balanced workload. In an embodiment, each computer receives a respective initial partition that contains a respective subset of edges of a graph and a respective subset of vertices the graph. For each edge of its respective subset of edges, each computer stores an intermediate representation of the edge into an edge table that resides on the computer. For each vertex of its respective subset of vertices, each computer stores an intermediate representation of the vertex into a vertex table that resides on the computer. The intermediate representation is optimized for mutability, which allows vertices and edges to be readily reallocated to other computers for load balancing.

(16) Different edges of a same vertex may be initially loaded by different computers. Thus, a vertex may be redundantly loaded on multiple computers, which may cause a potential conflict. To facilitate deconfliction, the computer announces that the vertex resides on the computer to a respective tracking computer of the set of computers. Each tracking computer makes assignments of particular vertices to particular computers and makes those assignments available to the set of computers. Each computer of the set of computers that initially loaded some of the particular vertices transfers the particular vertices and edges that originate at the particular vertices to the particular computers of the respective assignments. Each computer of the set of computers stores a materialized representation of a respective balanced partition based on: the edge table that resides on the computer, the vertex table that resides on the computer, and the vertices and the edges that

were transferred to the computer. Edges stored in the materialized representation of the respective balanced partition are stored in a different format than edges stored in the edge table that resides on the computer.

(17) In embodiments, a segmented array is used to implement each column of an edge table or vertex table. A segmented array provides mutability and random access with desirable asymptotic computational complexity. A technique for merging segmented arrays is presented herein.

(18) In embodiments, a table view is superimposed upon an edge table or vertex table to achieve thread safety and information hiding. In embodiments, batching and iteration are used for deconfliction of redundantly loaded vertices. In embodiments, ghost vertices are synthesized to avoid conflicts. In embodiments, coordination of the topology of computers is enhanced with a master computer that can orchestrate rebalancing of partitions.

2.0 Example Distributed System

(19) FIG. 1 is a block diagram that depicts an example distributed system **100**, in an embodiment. Distributed system **100** uses minimal communication to repartition a distributed graph to achieve a balanced workload.

2.1 Physical Topology

(20) Distributed system **100** contains networked computers such as **111-112** and **120**, which may be personal computers, rack servers such as blades, mainframes, virtual machines, or other connected computing devices capable of graph processing. Computers **111-112** and **120** intercommunicate over connections that span one or more networks such as a local area network (LAN), backplane, or internetwork.

2.2 Initial Partition

(21) To process distributed graph **130** in parallel, computers **111-112** respectively load initial partitions **141-142** of graph **130** from a file, database, or other shared data store. Graph **130** is a directed graph that has vertices **161-166** interconnected by directed edges A-F.

(22) Initial partitions **141-142** may contain more or less equal amounts of vertices of graph **130**. For example, initial partitions **141-142** may contain a same amount of vertices or edges or vertices and edges loaded from a same amount of bytes, text, or lines of text.

(23) As shown, initial partition **141** contains vertices **161-163**. Initial partition **142** contains a same amount of vertices, **164-165**.

(24) Just because initial partitions **141-142** contain a same amount of vertices does not mean that computers **111-112** have balanced workloads. For example, initial partitions **141-142** may have unbalanced amounts of edges.

2.3 Graph Retrieval

(25) Which edges A-F occur in which initial partitions **141-142** depends on the embodiments. For example, Oracle PGX supports a variety of graph source file formats such as adjacency list, edge list, two tables text, and flat file that may serialize vertices and edges of a same graph in different orders.

(26) For example, graph **130** may be originally encoded in two table text format. For example, one text file may encode vertices **161-166** each having color and month properties with the following six lines of text. Line 1: **161** blue April Line 2: **162** red February Line 3: **163** blue June Line 4: **164** green April Line 5: **165** red June Line 6: **166** yellow May

(27) Likewise, another text file may encode edges A-F each having a weight property with the following six lines of text. For example, Line 1 encodes edge A as originating from vertex **162**, terminating at vertex **164**, and having a weight of 1.5. Line 1: **162 164 1.5** Line 2: **166 162 1.1** Line 3: **165 166 3.0** Line 4: **165 164 0.4** Line 5: **161 164 0.4** Line 6: **165 163 1.1**

(28) Initial partition **141** may be populated from lines 1-3 of the vertex text file and lines 1-3 of the edges text file. Thus, initial partition **141** may contain vertices **161-163** and edges A-C. Likewise, initial partition **112** may contain vertices **164-166** and edges D-F.

(29) Thus with the two tables text format, initial partitions **111-112** may contain an equal amount of

vertices and may contain another (possibly different) equal amount of edges. However, initial partitions **141-142** have poor topological locality of data.

(30) Thus, computers **111-112** may have unbalanced communication and/or processing loads. For example, initial partition **141** contains edge C that connects vertices **165-166** that reside in a different initial partition **142**.

(31) Other graph source formats may cause unbalanced partitioning/loading in different ways. For example if a directed edge is initially loaded with the vertex from which the edge originates, then initial partition **141** may contain vertices **161-163** but contain only edges A and E, which are the only two edges that originate from vertices **161-163**.

(32) For example, graph **130** may be originally encoded in adjacency list text format. For example, a sole text file may encode vertices **161-166** and edges A-F with the following six lines of text (one line per vertex). Line 1: **161** blue April **164** 0.4 Line 2: **162** red February **164** 1.5 Line 3: **163** blue June Line 4: **164** green April Line 5: **165** red June **166** 3.0 **164** 0.4 **163** 1.1 Line 6: **166** yellow May **162** 1.1

(33) Thus with the above adjacency list, initial partition **141** contains only edges A-B, shown in bold in text lines 1-2. Whereas, initial partition **142** may have twice as many edges (C-F), shown in bold in text lines 4-6, despite having a same amount of vertices.

2.4 Vertex Table and Edge Table

(34) Regardless of which graph source format, computer **111** loads initial partition **141** into local volatile or non-volatile storage. Computer **111** may load initial partition **141** into data structures such as tables **170** and **180**.

(35) Computer **111** encodes vertices **161-162** as respective intermediate vertex representations **181-183** in vertex table **180**. Likewise, computer **111** encodes edges A-B as respective intermediate edge representations **171-172** in edge table **170**. Likewise, computer **112** loads initial partition **142**.

2.5 Tracking Computer

(36) Each of vertices **161-166** is associated with a tracking computer, such as **120**, that records which computer owns (e.g. loaded) which vertices. Although not shown, tracking computer **120** may load a respective initial partition. Likewise, computers **111-112** may also operate as tracking computers.

(37) In embodiments, each of computers **111-112** and **120** loads a partition and tracks vertices. Which vertices are tracked by tracking computer **120** may depend on the implementation.

(38) Regardless of the embodiment, all vertices are tracked by exactly one computer. Furthermore, each vertex is tracked by a same computer for the duration of the balancing process described herein.

(39) Computers **111-112** detect which computers track which vertices of respective initial partition **141-142**. Computers **111-112** announce to tracking computer **120** which vertices (and how many edges of those vertices) have computers **111-112** loaded.

(40) For example computer **112** may send a network message, such as a hypertext transfer protocol (HTTP) request, that announces to tracking computer **120** that computer **112** loaded vertex **165** and two (D and F) of the edges of vertex **165**. Likewise, computer **112** may announce to a different tracking computer that computer **112** loaded vertex **166** without any vertices, even though vertex **166** actually has two edges (that are loaded by computer **111**). Likewise, computer **111** may announce to tracking computer **120** that computer **111** loaded edge C that originates from vertex **165**, even though vertex **165** is loaded and announced by computer **112**.

2.6 Vertex Reassignment

(41) In operation, tracking computers record assignments that indicate which computers own which vertices. For example, tracking computer **120** contains assignment **125** that indicates that vertex **165** is owned by computer **112**.

(42) Assignments may initially reflect which computer initially loaded which vertex. However, a conflict arises when multiple computers report a same vertex (to a same tracking computer).

- (43) For example, edges C-D and F of vertex **165** are loaded by different computers. Computer **111** loads edges D and F. Whereas, computer **112** loads edge C.
- (44) Thus, computers **111-112** both announce vertex **165** to tracking computer **120**. However, the announcement by computer **112** may indicate that computer **112** actually loaded the properties of vertex **165**. Whereas, computer **111** merely incidentally loaded vertex **165**, without properties, as an endpoint of an edge that computer **111** actually loaded.
- (45) For example, vertex **165** may have properties for product type, product number, quantity, order date, delivery date, delivery address that are loaded by computer **112** and not computer **111**. Only the vertex table of computer **112** may contain a row for the intermediate representation of vertex **165**.
- (46) Assignments are mutable, and tracking computers may update the assignments that they maintain. For example, tracking computer **120** may direct computers **111-112** to reassign ownership of vertices (and their edges) to other computers. Computers **111-112** process such reassignment directions by annotating intermediate representations **171-172** and **181-183** to record which vertices and edges should eventually be moved to which other computers.
- (47) Each row of tables **170** and **180** may record a respective edge or vertex. Tables **170** and **180** are mutable and may dynamically store additional rows.
- (48) For example, tracking computer **120** may inform computers **111-112** that vertex **164** is reassigned from computer **112** to computer **111**. Thus, computer **112** may mark its intermediate vertex representation (not shown) of vertex **164** to indicate that vertex **164** is now owned by computer **111**.
- (49) However to accomplish that reassignment of vertex **164**, computers **111-112** need not communicate with each other, because tracking computer **120** mediates the reassignment. Furthermore, computer **112** may have initially loaded many properties (not shown) for vertex **164**.
- (50) Thus, computer **112** may have an intermediate vertex representation of vertex **164** that has many properties and indicates that vertex **164** is reassigned to computer **111**.
- (51) Thus, reassignment of a vertex (along with its edges that originate from a source computer) may occur at any time as directed by a tracking computer from the source computer to a target computer merely by annotating the vertex and its edges as moved in the vertex table and edge table of the source computer without transferring properties of the vertex and edges. Thus, reassignment of vertices and edges is inexpensive in terms of processing and communicating.
- (52) To reassign a vertex, only the vertex's identifier need be sent. An identifier may be a primitive integer such as a serial number, text line number, or database table row number. To virtually move an edge, only the identifiers of the originating vertex and terminating vertex need be sent.
- (53) Thus, vertices and edges may be batched at very high density for reassignment. Thus, large amounts of vertices and edges may be reassigned.
- (54) Likewise, a same vertex and edges may be repeatedly virtually moved without significant cost. Thus, repeated movement and batching naturally support virtual movement of vertices and edges in iterations and in phases having iterations.
- (55) As explained later herein, computers **111-112** may have limited intercommunication with tracking computers during the iterations so that distributed system **100** may eventually converge on a near optimal assignment of vertices to computers. Optimal assignment maximizes edge locality such that a vertex and edges it originates are assigned to a same computer.

2.7 Virtual Movement

- (56) Virtual movement is even cheaper than reassignment. Virtual movement entails a tracking computer updating its tracking data to indicate that a vertex has moved from one computer to another, without immediately notifying the computers about the movement. Virtual movement may occur over multiple iterations, and a same vertex can be repeatedly virtually moved without any network communication until a next iteration or after all iterations.

2.8 Actual Movement

(57) Eventually the tracking computers or a master computer (not shown) of distributed system **100** detects convergence or a maximum iterations and may broadcast to computers **111-112** that reassignments now need execution. Responsively, computers **111-112** may send the property values of vertices (and edges) that computers **111-112** initially loaded to whichever computers those vertices and edges are now assigned. For example, computer **111** may send to computer **112** some or all of the data of a row of table **170** or **180** to accomplish a transfer of ownership of an edge or vertex.

(58) Thus, reassignments may describe a balanced data distribution, and subsequent transfer of properties or whole table rows may finally achieve the described balance. Henceforth, computers **111-112** may analyze or otherwise process graph **130** in a distributed and balanced way.

(59) However, tables **170** and **180** are designed for mutation. Whereas, graph analytics may treat graph **130** as more or less immutable.

(60) Thus prior to actual analysis and for efficiency, computers **111-112** should recast their loaded graph data into other data structures that are optimized for read-only use, such as materialized representation **150**. For example, computer **111** uses the contents of tables **170** and **180** to populate materialized representation **150**. Implementations of mutable and immutable data structures are discussed later herein.

3.0 Example Repartitioning Process

(61) FIG. **2** is a scenario diagram that depicts example distributed system **100** in operation. Thus, FIG. **2** is discussed with reference to FIG. **1**.

(62) FIG. **2** shows interactions as horizontal arrows between components. For example at time A, computer **111** receives initial partition **141**. Although not shown, computer **112** receives initial partition **142** more or less at time A also.

(63) For example, computers **111-112** may each respectively load a contiguously stored portion of graph **130**, such as an equal amount of rows from a database table or lines of text from a file. Alternatively, graph **130** may be loaded partially or totally from a stream or multiple streams of data, such as time series data, such as a stock exchange ticker tape or concurrent ticker tapes from rival exchanges. No matter what data store(s) provides the details of graph **130**, distributed system **100** needs only one pass over the data of graph **130** to achieve a distributed loading of graph **130**.

(64) At times B-C, which may be a same time, computers **111-112** store encode their initial partitions into mutable local data structures. For example, computer **111** encodes each vertex or edge of initial partition **141** as an intermediate representation row in table **170** or **180**. Thus immediately after time C, graph **130** is loaded, although likely unbalanced.

(65) At time D, computers **111-112** detect which computers track which vertices. Also at time D, computers **111-112** inform those tracking computers that various vertices are loaded along with a count of edges that those vertices originate.

(66) At time E and based on the loading information provided by computers **111-112**, tracking computers such as **120** decide which vertices should be reassigned to which other computers. Tracking computer **120** reassigns vertices to eliminate situations where multiple machines store different edges for a same vertex.

(67) Thus at time E, tracking computer **120** updates assignments, such as **125**, and then provides the updated assignments to computers that gain or lose ownership of involved vertices. Edge counts may be involved to decide reassignments, and edges are themselves reassigned when a vertex is reassigned.

(68) As explained later herein and although not shown, the interaction at time D may be iteratively repeated. Subsequent announcements at time D may be limited to conflicted vertices (those with edges loaded by different computers).

(69) Thus, information sent at time E may include identification of vertices that remain conflicted. As the process iteratively converges, the amounts of vertices involved at times D-E diminish. The amounts of vertices involved at times D-E may also be limited by batching, such that not all

conflicted vertices are addressed in a same iteration, as explained later herein.

(70) As explained above, initial partitioning may have unequal amounts of edges or vertices. Likewise, deconfliction by tracking computers may introduce or aggravate such an imbalance.

(71) Thus at time F, rebalancing occurs. As explained later herein and although time F shows only a single interaction, rebalancing may actually be iterative and involve supervision by a master computer (not shown). However like reassignment, rebalancing merely entails annotating intermediate representations of vertices and/or edges as moved and updating assignments by tracking computers rather than actually moving the properties of vertices and edges.

(72) As a more or less final activity at time F, after all reassignment and rebalancing is done, vertices and edges (including their identifiers and properties) are actually sent to the computers designated to finally host them. For example, a vertex may be virtually moved during deconfliction and/or virtually moved during rebalancing, perhaps repeatedly to various computers.

(73) However, ultimately a tracking computer recorded some computer as finally owning a given vertex. Time F more or less finishes with computers forwarding vertices and edges (including their identifiers and properties) to other computers as indicated by the tracking computers.

(74) Thus while times D-F may virtually shuffle vertices around as an accounting effort, time F finishes by actually moving the vertices to computers to achieve an actually balanced distribution. Such actual movement of vertices and edges may be done in bulk. Also at time F after the actual movement, computers **111-112** populate their materialized representations, such as **150**, to reflect the final partitioning in a format that is optimized for reading during actual graph analysis. For example, a compressed sparse row format or a compressed sparse column format may be used as the materialized format, such as described in U.S. Pat. No. 8,775,495, entitled "Compression system and method for accelerating sparse matrix computations." A more detailed discussion of activity during time F is provided later herein.

4.0 Intermediate Representation

(75) FIG. 3 is a block diagram that depicts an example computer **300**, in an embodiment. Computer **300** uses a segmented array to implement each column of an edge table or vertex table.

(76) Computer **300** may be an implementation of computer **111**. Computer **300** stores tables **370** and **380** in volatile memory.

4.1 Segmented Array

(77) A segmented array is used for each column of tables **370** and **380** to optimize the performance of mutation. Thus, segmented arrays **311-313** are designed to efficiently append additional rows.

(78) Thus after rebalancing, edges and rows that are actually moved to computer **300** from another computer may be rapidly inserted into tables **370** and **380**. To facilitate rapid growth, segmented arrays **311-313** have a hybrid data structure that combines aspects of conventional arrays and linked lists.

(79) A conventional array, such as a vector in the C++ standard template library, is prone to unpredictable latency during growth because growth entails copying all existing data in addition to allocating logarithmically more blank rows than actually needed. Thus, a conventional array has unstable growth performance.

(80) Whereas, a linked list grows in constant time. However, a linked list may fragment a heap and thrash a CPU cache.

(81) By using arrays as nodes (segments) of a linked list, a segmented array incrementally grows in constant time. A segment may be sized for a cache line. When segments share a fixed size and are indexed (not shown), random access of a row in a table column occurs in constant time when the column is implemented as a segmented array.

(82) Edge table **370** topologically encodes edges as rows in a table that has source and destination columns to identify pairs of vertices that directed edges connect. For example, edge E3 originates at vertex V2 and terminates at vertex V3.

(83) The row and edge columns of edge table **370** are for demonstrative purposes and are implied.

Thus, the row and edge columns may be absent from some embodiments.

(84) The row column of segments **321-322** also is demonstrative and implied. In embodiments, segments **321-322** are merely sequences of values stored contiguously, shown as a value column.

(85) Vertices and edges may have many properties, with each property having its own table column implemented with its own segmented array. Actual movement of an edge requires appending a row into edge table **370**, because various vertices and edges need materialization by actually transferring their properties. That transfer may include properties not previously encountered by a given computer.

(86) For example, computer **300** may load an initial partition that contains edges E1-E3. As shown, edges E1-E3 have no custom properties.

(87) Whereas, edge E99 and its label property may finally move to computer **300**. Even though edge table **370** might not already have a label column, a label column and its segmented array may be created and added to edge table **370** just in time. Thus, tables **370** and **380** may lazily add (or remove) columns when needed.

(88) For example, a column may be dynamically added to annotate a table with ownership information or dynamically removed to discard metadata (such as ownership information) no longer needed. However, actual identifier and property data are loaded and transferred together.

5.0 Segments

(89) FIG. **4** is a block diagram that depicts an example computer **400**, in an embodiment. Computer **400** optimally merges two segmented arrays.

(90) Computer **400** may be an implementation of computer **111**. Segmented arrays **411-412** are stored in the memory of computer **400**.

5.1 Random Access

(91) The segments of each segmented array **411-412** are indexed by a lookup vector of pointers. For example, segmented array **411** has two segments, **421A** and **422**. Thus, lookup vector **431** has two pointers that point to respective segments **421A** and **422**.

(92) Thus, random access of a particular segment occurs in constant time. When segments **421A** and **422** have a same capacity (as shown), random access of any value in segmented array **411** occurs in constant time.

(93) In the shown embodiment, each segment has a next pointer to a next segment. Thus, segment **421A** has a next pointer to segment **422**. Segment **422** may have a null value for its next pointer because segment **422** is the last segment of segmented array **411**. In embodiments not shown, segments **421-422** do not have next pointers, and identification of a next segment entails dereferencing another pointer in a lookup vector such as **431**. For example, the next pointer from full segments **421A** to partial segment **422** may be absent and instead approximated by the pointer to partial segment **422** in lookup vector **431**. Thus, computer **400** may instead iterate the segments of segmented array **411** by dereferencing successive contiguous pointers that occur in lookup vector **431** as shown.

5.2 Optimal Merge

(94) Segmented arrays **411-412** may store values of a same table column (not shown). For example, actual movement of multiple vertices and/or edges may entail transmission of reassigned rows of a vertex table or edge table from one computer to another. Upon receipt of actually moved vertices or edges, the receiving computer may briefly store the received rows in temporary segmented arrays. To fulfil the actual movement of vertices or edges, the temporary segmented arrays may need to be merged into the vertex table and edge table of the receiving computer. For example, segmented array **411** may be temporary and for a same vertex column or edge column as segmented array **412**. Thus eventually, computer **400** may need to merge segmented array **411** into segmented array **412A**, which is shown as segmented array **412A** becoming segmented array **412B**. Such a merge mutates segmented array **412** and consumes (destroys) segmented array **411**.

(95) Computer **400** is optimized to minimize data copying when merging segmented arrays. Thus,

computer **400** merges segmented arrays in optimal time.

(96) Computer **400** accomplishes an optimal merge primarily by adjusting segment pointers and secondarily by copying at most one partial segment such as **422**. A partial segment is not fully populated. As shown, segmented arrays **411** and **412A** may each have one partial segment as a last segment.

(97) Thus, segmented array **411** contains at most one partial segment (**422**) and at least one full segment (**421A**). During a merge, computer **400** reassigns the next pointer of the last full segment (**421A**) of segmented array **411** to point to the first segment of segmented array **412A**.

(98) Thus, the first segment of segmented array **412B** is shown as segment **421B**, which actually was segment **421A**. Thus, computer **400** prepends all of the full segments of segmented array **411** onto the beginning of segmented array **412A**.

(99) Computer **400** accomplishes that by reassigning one next pointer and without copying any column values. Computer **400** also prepends pointers to those inserted segments (e.g. **421B**) onto the beginning of lookup vector **432B**. Such prepending occurs in time that is linear with the total number of segments of segmented arrays **411** and **412A**, which may be orders of magnitude less than the number of rows in the edge table.

(100) Computer **400** also appends the contents of partial segment **422** into partial segment **424A**, shown as combined segment **424B** that was partial segment **424A**. Such appending occurs in time that is linear with the fixed size that all segments have, which may be orders of magnitude less than the number of rows in the edge table. If necessary (not shown), computer **400** may reuse partial segment **422** by adjusting the next pointer (and lookup vector **432B**) to point to partial segment **422**, such as when the merge causes combined segment **424B** to overflow. Thus as described, computer **400** may merge two segmented arrays in time that is less (better) than linear with the number of rows in the edge table.

6.0 Table Optimizations

(101) FIG. 5 is a block diagram that depicts an example computer **500**, in an embodiment. Computer **500** has various optimizations that accelerate the use of a vertex table or edge table. Computer **500** may be an implementation of computer **111**.

6.1 Indexing

(102) The rows of edge table **570** may occur in a same ordering as initially loaded, which as explained above, may vary according to implementation details such as source file format. Thus, the rows of edge table **570** may be arbitrarily ordered in a way that does not reveal an ordering that is natural to the edge property data.

(103) For example, the values in the month column of edge table **570** are not in calendar order. For example, August does not naturally occur between March and May, even though edge table **570** shows such an unnatural ordering of months.

(104) Index **530** may provide fast access to the months (and the edges themselves) of edge table **570** in a sorted ordering, such as a natural order, such as according to a calendar. The row numbers that are stored in index **530** are sorted by month according to the calendar. Thus, row 1 is listed last in index **530**, because in row 1 of edge table **570** is August, which is the naturally last of the values in the month column when calendar sorted.

(105) Indexing may serve two purposes, one of which is sorting as explained above. The other purpose is aggregation.

(106) For example, enumeration of neighbor vertices of a current vertex may involve selecting all edges that originate at the current vertex. Thus, the source column of edge table **570** may have its own index that merely groups edges by source vertex without actual numeric sorting.

6.2 Multithreading

(107) Computer **500** has execution threads **541-542** that may run concurrently when allocated to separate CPUs or cores, such as with symmetric multiprocessing (SMP).

(108) All of the edges owned (initially loaded or transferred from other computers) by computer

500 are stored as rows in edge table **570**. Edge table **570** provides a separate table view of edge table **570** to each of threads **541-542**.

(109) Table views share a segmented array for each column, such as segmented array **510** for the month column. Edge table **570** is inherently thread safe for concurrent reading.

(110) In embodiments, some operations (e.g. row deletion) are not inherently thread safe, and synchronization may be needed to achieve thread safety.

6.3 Table View

(111) The horizontal bold dashed line shows that table view for thread **541** exposes the rows of the top half of edge table **570**. Likewise, the table view for thread **542** exposes the rows of the bottom half of edge table **570**.

(112) However, such as neat split is not essential. For example, an embodiment may instead expose every odd row through the table view of thread **541** and expose every even row through the table view of thread **542**. Indeed, insertion of an additional row may occur by appending to the bottom of edge table **570** regardless of which thread should view the additional row.

(113) As described, a table view may project a subset of rows of an underlying table. Likewise by including a subset of segmented arrays of the underlying table, a table view may project a subset of columns.

(114) Growth (addition of rows) of edge table **570** may be inherently thread safe so long as each table view exposes a non-overlapping and sufficiently pre-allocated portion (e.g. whole segments) of segmented array **510**. For example, segments **521-522** are allocated for thread **542**, and segments **523-524** are for thread **541**. Segment **522** has spare capacity for an additional value to be inserted later into thread **542**'s table view.

(115) The table view of each thread has its own lookup vector that contains pointers to indicate which segments belong to which thread. For example, lookup vector **532** has pointers that indicate that segments **521-522** belong to thread **542**.

(116) Non-overlapping, as mentioned above, means that no row of edge table **570** is shared by multiple table views. However, rows owned by different table views may be interleaved (not shown) within edge table **570**.

(117) As shown, each of lookup vectors **531-532** has pointers to adjacent segments. For example, all of thread **542**'s segments **521-522** occur in segmented array **510** before all of thread **541**'s segments **523-524**. However, the segments of threads **541-542** may be interleaved (not shown), such that a segment for thread **541** occurs between two segments for thread **542**.

7.0 Logical Topology

(118) The techniques described above may be combined in various ways to load and adjust graph partitions. FIG. **6** shows a computer cluster arranged in a logical topology that is dynamically managed by tracking computers. FIG. **6** introduces a top tier consisting of a sole master computer that imposes central control to direct the tracking computers. FIGS. **7-9** shows additional techniques, more or less based on that cluster, for further coordination.

(119) FIG. **6** is a block diagram that depicts an example cluster **600**, in an embodiment. Cluster **600** has a central master computer to manage tracking computers.

(120) Cluster **600** may be an implementation of distributed system **100**. Cluster **600** contains computers **611-613**, **621-622**, and **680**.

(121) Computers **611-613** each load a partition of a graph. Some vertices of the graph, such as **673-676**, may be redundantly loaded by two or more of computers **611-613**.

(122) For example, computer **611** may load vertex **673** as part of loading edges A-B.

(123) Likewise, computer **612** may also (redundantly) load vertex **673** as part of loading edge C. Redundant loading may need deconfliction as detected, decided, and directed by tracking computers **621-622**.

(124) Tracking computers **621-622** track which computers own which vertices and how many edges and may direct computers **611-613** to redistribute vertices and edges in various ways.

Tracking computers **621-622** may centrally coordinate through master computer **680**.

(125) Computers **621-622** and **680** may each also load a respective partition. Master computer **680** may also be a tracking computer. Cluster **600** has only one master computer.

(126) Additional coordination techniques that use cluster **600**, with or without master computer **680**, are presented with FIGS. 7-9 and discussed as follows.

8.0 Deconfliction

(127) FIG. 7 is a scenario diagram that depicts an example technique for operating cluster **600**, in an embodiment. This technique uses batching and iteration for deconfliction of redundantly loaded vertices.

(128) This technique includes a preparatory phase and subsequent iterations, such as **731-732**. The preparatory phase occurs at time A.

(129) At time A computers **611-613** each loads a respective partition of a graph. Also at time A, computers **611-613** hash an identifier of each vertex they load to detect which of their vertices are tracked by which of tracking computers **621-622**.

(130) For example, cluster **600** may identify each of tracking computers **621-622** by a zero-based serial number and each of the vertices of the entire graph by a zero-based ordinal (offset) into a global list of vertices. A hash function may take a vertex ordinal as an argument and emit a tracking computer serial number.

(131) For example, the hash function may divide the vertex ordinal by a count of tracking computers of cluster **600** and return the remainder (modulo) as a tracking computer serial number. Thus, cluster **600** may identify a tracking computer for each vertex of the graph, and each tracking computer tracks a same amount of vertices.

(132) Thus, computers **611-613** may each load vertices and detect which of tracking computers **621-622** tracks each loaded vertex. Thus, computers **611-613** may detect which of their vertices are tracked by tracking computer **621**.

(133) Also at time A, computers **611-613** announce to tracking computers **621-622** which vertices are loaded. At time B and based on the announcements of loaded vertices, tracking computers **621-622** may detect vertex conflicts.

(134) For example, tracking computer **621** may detect that vertices **673-676** are redundantly loaded by computers **611-613**. However, only redundant vertices that originate edges from different computers are conflicted. In embodiments at time A, computers **611-613** announce to tracking computers **621-622** which vertices are redundantly loaded and/or which vertices are conflicted.

(135) For example, vertex **677** is not conflicted even if computer **611** loaded edge C because vertex **677** does not originate any edges that terminate at a redundantly loaded vertex. For example even if computer **612** loaded vertex **677**, another vertex (not shown), and an edge that connects those two vertices and originates at vertex **677**, then vertex **677** would still not be conflicted because the edge and its endpoints are local

(136) Also at time B in embodiments (not shown), tracking computers **621-622** may announce all of the vertices of the graph that are conflicted to computers **611-613**. However, an attempt to exhaustively deconflict in a single pass may need an amount of memory on the order of the multiplicative product of the number of vertices of the graph times the number of computers **611-613** in cluster **600**.

(137) In the shown embodiment, tracking computers **621-622** instead announce conflicted vertices in batches having a same fixed amount of vertex identifiers. In the shown embodiment, tracking computer prepares a first batch of vertex identifiers from computer **611** that are conflicted, a second batch of vertex identifiers from computer **612** that are conflicted, and a third batch (not shown) of vertex identifiers from computer **613** that are conflicted.

(138) Thus in the shown embodiment, the three batches identify disjoint (non-overlapping) or partially overlapping sets of vertices. At time B, tracking computer **621** sends each of those batches to a respective computer of **611-613**. In embodiments (not shown), tracking computer **621** prepares

only one batch instead of three and broadcasts it to all computers **611-613**.

(139) At time C, each of computers **611-613** respond to a received conflict batch by reporting to tracking computer **621** how many edges are currently owned by the computer. For example, computer **611** may report for each conflicted vertex of the batch a count of that vertex's inbound edges currently owned by computer **611** and a count of that vertex's outbound edges currently owned.

(140) After iteration **731**, the process generally repeats the process during iteration **732**. Iteration continues until no vertices are conflicted.

(141) At time F, tracking computers **621-622** process the received edge counts and reassigns the respective conflicted vertices to one of computers **611-613** based on which computer already owns the most edges that are connected to a respective conflicted vertex. For example, tracking computer **621** may designate computer **611** as sole owner of vertex **673**, because computer **611** has more edges (A-B) for vertex **673** than does computer **612**.

(142) If a vertex originates no edges, tracking computer **621** reassigns the vertex to whichever computer initially loaded the properties of the vertex. Edges of vertices that are not conflicted are reassigned to whichever computer owns the vertex that originates the edge.

(143) Also at time F, tracking computer **621** moves any vertices and edges that it reassigns by sending reassignments to computers **611-613**.

9.0 Ghost Vertex

(144) FIG. **8** is a scenario diagram that depicts an example technique for operating cluster **600**, in an embodiment. This technique uses ghost vertices to avoid conflicts.

(145) This technique is performed before first iteration **731** of deconfliction. Indeed, times A-C of FIG. **8** occur during time A of FIG. **7**.

(146) The more edges a vertex has, the more likely the vertex is to be redundantly loaded (and thus potentially conflicted). Furthermore even after deconfliction, a vertex of high degree (edge count) may still have many edges that originate or terminate at vertices owned by other computers.

(147) Thus, vertices of high degree may cause excessive communication during subsequent graph analytics, even after repartitioning is done. A ghost vertex is a lightweight proxy that represents a same vertex on all computers, although only one computer actually loads the properties of the ghost vertex.

(148) Thus, subsequent analysis of a ghost vertex and a subset of its edges may occur entirely local to any of the various computers and without network communication. Thus, a ghost vertex is not subjected to deconfliction.

(149) At time A, tracking computers **621-622** detect which of their tracked vertices have a degree that exceeds a threshold. Those are the high degree vertices.

(150) Also at time A, tracking computers **621-622** rank (sort) the high degree vertices by degree and select at most a first fixed amount of those vertices having the highest degrees. Those are the ghost candidate vertices. Also at time A, tracking computers **621-622** report an identifier and edge count for each candidate vertex to master computer **680**.

(151) At time B, master computer **680** ranks the candidate vertices by their degree and promotes at most a second fixed amount of them to actually be ghost vertices. Also at time B, master computer **680** announces the identifiers of the ghost vertices to computers **611-613**.

(152) Each of computers **611-613** has a separate data structure, such as ghost set **870**, to record which vertices are ghosts. Ghost set **870** is separate and in addition to the vertex table of each computer. At time C, each of computers **611-613** insert all of the ghost vertices into its own ghost set.

(153) After all phases and iterations (e.g. FIG. **7**) of virtual movement of vertices and edges are done, cluster **600** actually repartitions by actually moving vertices and edges (including their identifiers and properties) to their decided owner computers. During such actual movement, a computer that initially loaded properties of a moving vertex or edge actually sends those properties

to another computer.

(154) During such actual movement, edges outgoing/incoming from/to a ghost vertex are moved to whichever computer owns the non-ghost end of an edge. Edges between two ghost vertices are not moved.

10.0 Rebalancing

(155) FIG. **9** is a scenario diagram that depicts an example technique for operating cluster **600**, in an embodiment. This technique uses the master computer to rebalance partitions.

(156) This technique is iteratively performed, although FIG. **9** shows only one iteration. In embodiments, one or more iterations of this technique occur after each of deconfliction iterations **731-732** of FIG. **7**. In embodiments, this technique performs its iterations only after all deconfliction iterations have finished.

(157) At time A, tracking computers **621-622** each report to master computer **680** how many edges originate and how many edges terminate at a computer for vertices that are tracked by the tracking computer.

(158) At time B based on those counts, master computer **680** detects which computers have too many edges and thus are imbalanced. Master computer **680** may invoke various rules, heuristics, and logic to decide how many edges should some computers virtually move and to which other computers.

(159) Also at time B based on that deciding, master computer directs each tracking computer to transfer a respective first count of edges from a first computer to a second computer, and a second count of edges to transfer from a third computer to a fourth computer, and so forth. At time C and based on the respective transfer counts, each of tracking computers **621-622** may invoke various rules, heuristics, and logic to decide which surplus vertices should be virtually moved to approximately achieve the edge transfer counts.

(160) Further processing at time C depends on the embodiment. In embodiments that perform this technique only after finishing all deconfliction iterations, processing at time C occurs in between the processing of times D-E of FIG. **2**.

(161) In embodiments, processing at time C may be included in processing at time D of FIG. **7**. Regardless of embodiment at time C, master computer **680** directs tracking computers **621-622** to virtually move which surplus vertices (and their edges) to which other computers.

11.0 Hardware Overview

(162) According to one embodiment, the techniques described herein are implemented by one or more special-purpose computing devices. The special-purpose computing devices may be hard-wired to perform the techniques, or may include digital electronic devices such as one or more application-specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs) that are persistently programmed to perform the techniques, or may include one or more general purpose hardware processors programmed to perform the techniques pursuant to program instructions in firmware, memory, other storage, or a combination. Such special-purpose computing devices may also combine custom hard-wired logic, ASICs, or FPGAs with custom programming to accomplish the techniques. The special-purpose computing devices may be desktop computer systems, portable computer systems, handheld devices, networking devices or any other device that incorporates hard-wired and/or program logic to implement the techniques.

(163) For example, FIG. **10** is a block diagram that illustrates a computer system **1000** upon which an embodiment of the invention may be implemented. Computer system **1000** includes a bus **1002** or other communication mechanism for communicating information, and a hardware processor **1004** coupled with bus **1002** for processing information. Hardware processor **1004** may be, for example, a general purpose microprocessor.

(164) Computer system **1000** also includes a main memory **1006**, such as a random access memory (RAM) or other dynamic storage device, coupled to bus **1002** for storing information and instructions to be executed by processor **1004**. Main memory **1006** also may be used for storing

temporary variables or other intermediate information during execution of instructions to be executed by processor **1004**. Such instructions, when stored in non-transitory storage media accessible to processor **1004**, render computer system **1000** into a special-purpose machine that is customized to perform the operations specified in the instructions.

(165) Computer system **1000** further includes a read only memory (ROM) **1008** or other static storage device coupled to bus **1002** for storing static information and instructions for processor **1004**. A storage device **106**, such as a magnetic disk or optical disk, is provided and coupled to bus **1002** for storing information and instructions.

(166) Computer system **1000** may be coupled via bus **1002** to a display **1012**, such as a cathode ray tube (CRT), for displaying information to a computer user. An input device **1014**, including alphanumeric and other keys, is coupled to bus **1002** for communicating information and command selections to processor **1004**. Another type of user input device is cursor control **1016**, such as a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **1004** and for controlling cursor movement on display **1012**. This input device typically has two degrees of freedom in two axes, a first axis (e.g., x) and a second axis (e.g., y), that allows the device to specify positions in a plane.

(167) Computer system **1000** may implement the techniques described herein using customized hard-wired logic, one or more ASICs or FPGAs, firmware and/or program logic which in combination with the computer system causes or programs computer system **1000** to be a special-purpose machine. According to one embodiment, the techniques herein are performed by computer system **1000** in response to processor **1004** executing one or more sequences of one or more instructions contained in main memory **1006**. Such instructions may be read into main memory **1006** from another storage medium, such as storage device **106**. Execution of the sequences of instructions contained in main memory **1006** causes processor **1004** to perform the process steps described herein. In alternative embodiments, hard-wired circuitry may be used in place of or in combination with software instructions.

(168) The term “storage media” as used herein refers to any non-transitory media that store data and/or instructions that cause a machine to operation in a specific fashion. Such storage media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage device **106**. Volatile media includes dynamic memory, such as main memory **1006**. Common forms of storage media include, for example, a floppy disk, a flexible disk, hard disk, solid state drive, magnetic tape, or any other magnetic data storage medium, a CD-ROM, any other optical data storage medium, any physical medium with patterns of holes, a RAM, a PROM, and EPROM, a FLASH-EPROM, NVRAM, any other memory chip or cartridge.

(169) Storage media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between storage media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise bus **1002**. Transmission media can also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(170) Various forms of media may be involved in carrying one or more sequences of one or more instructions to processor **1004** for execution. For example, the instructions may initially be carried on a magnetic disk or solid state drive of a remote computer. The remote computer can load the instructions into its dynamic memory and send the instructions over a telephone line using a modem. A modem local to computer system **1000** can receive the data on the telephone line and use an infra-red transmitter to convert the data to an infra-red signal. An infra-red detector can receive the data carried in the infra-red signal and appropriate circuitry can place the data on bus **1002**. Bus **1002** carries the data to main memory **1006**, from which processor **1004** retrieves and executes the instructions. The instructions received by main memory **1006** may optionally be stored on storage device **106** either before or after execution by processor **1004**.

(171) Computer system **1000** also includes a communication interface **1018** coupled to bus **1002**. Communication interface **1018** provides a two-way data communication coupling to a network link **1020** that is connected to a local network **1022**. For example, communication interface **1018** may be an integrated services digital network (ISDN) card, cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of telephone line. As another example, communication interface **1018** may be a local area network (LAN) card to provide a data communication connection to a compatible LAN. Wireless links may also be implemented. In any such implementation, communication interface **1018** sends and receives electrical, electromagnetic or optical signals that carry digital data streams representing various types of information.

(172) Network link **1020** typically provides data communication through one or more networks to other data devices. For example, network link **1020** may provide a connection through local network **1022** to a host computer **1024** or to data equipment operated by an Internet Service Provider (ISP) **1026**. ISP **1026** in turn provides data communication services through the world wide packet data communication network now commonly referred to as the “Internet” **1028**. Local network **1022** and Internet **1028** both use electrical, electromagnetic or optical signals that carry digital data streams. The signals through the various networks and the signals on network link **1020** and through communication interface **1018**, which carry the digital data to and from computer system **1000**, are example forms of transmission media.

(173) Computer system **1000** can send messages and receive data, including program code, through the network(s), network link **1020** and communication interface **1018**. In the Internet example, a server **1030** might transmit a requested code for an application program through Internet **1028**, ISP **1026**, local network **1022** and communication interface **1018**.

(174) The received code may be executed by processor **1004** as it is received, and/or stored in storage device **106**, or other non-volatile storage for later execution.

12.0 Software Overview

(175) FIG. **11** is a block diagram of a basic software system **1100** that may be employed for controlling the operation of computing system **1000**. Software system **1100** and its components, including their connections, relationships, and functions, is meant to be exemplary only, and not meant to limit implementations of the example embodiment(s). Other software systems suitable for implementing the example embodiment(s) may have different components, including components with different connections, relationships, and functions.

(176) Software system **1100** is provided for directing the operation of computing system **1000**. Software system **1100**, which may be stored in system memory (RAM) **1006** and on fixed storage (e.g., hard disk or flash memory) **106**, includes a kernel or operating system (OS) **1110**.

(177) The OS **1110** manages low-level aspects of computer operation, including managing execution of processes, memory allocation, file input and output (I/O), and device I/O. One or more application programs, represented as **1102A**, **1102B**, **1102C** . . . **1102N**, may be “loaded” (e.g., transferred from fixed storage **106** into memory **1006**) for execution by the system **1100**. The applications or other software intended for use on computer system **1000** may also be stored as a set of downloadable computer-executable instructions, for example, for downloading and installation from an Internet location (e.g., a Web server, an app store, or other online service).

(178) Software system **1100** includes a graphical user interface (GUI) **1115**, for receiving user commands and data in a graphical (e.g., “point-and-click” or “touch gesture”) fashion. These inputs, in turn, may be acted upon by the system **1100** in accordance with instructions from operating system **1110** and/or application(s) **1102**. The GUI **1115** also serves to display the results of operation from the OS **1110** and application(s) **1102**, whereupon the user may supply additional inputs or terminate the session (e.g., log off).

(179) OS **1110** can execute directly on the bare hardware **1120** (e.g., processor(s) **1004**) of computer system **1000**. Alternatively, a hypervisor or virtual machine monitor (VMM) **1130** may be interposed between the bare hardware **1120** and the OS **1110**. In this configuration, VMM **1130**

acts as a software “cushion” or virtualization layer between the OS **1110** and the bare hardware **1120** of the computer system **1000**.

(180) VMM **1130** instantiates and runs one or more virtual machine instances (“guest machines”). Each guest machine comprises a “guest” operating system, such as OS **1110**, and one or more applications, such as application(s) **1102**, designed to execute on the guest operating system. The VMM **1130** presents the guest operating systems with a virtual operating platform and manages the execution of the guest operating systems.

(181) In some instances, the VMM **1130** may allow a guest operating system to run as if it is running on the bare hardware **1120** of computer system **1100** directly. In these instances, the same version of the guest operating system configured to execute on the bare hardware **1120** directly may also execute on VMM **1130** without modification or reconfiguration. In other words, VMM **1130** may provide full hardware and CPU virtualization to a guest operating system in some instances.

(182) In other instances, a guest operating system may be specially designed or configured to execute on VMM **1130** for efficiency. In these instances, the guest operating system is “aware” that it executes on a virtual machine monitor. In other words, VMM **1130** may provide para-virtualization to a guest operating system in some instances.

(183) A computer system process comprises an allotment of hardware processor time, and an allotment of memory (physical and/or virtual), the allotment of memory being for storing instructions executed by the hardware processor, for storing data generated by the hardware processor executing the instructions, and/or for storing the hardware processor state (e.g. content of registers) between allotments of the hardware processor time when the computer system process is not running. Computer system processes run under the control of an operating system, and may run under the control of other programs being executed on the computer system.

13.0 Cloud Computing

(184) The term “cloud computing” is generally used herein to describe a computing model which enables on-demand access to a shared pool of computing resources, such as computer networks, servers, software applications, and services, and which allows for rapid provisioning and release of resources with minimal management effort or service provider interaction.

(185) A cloud computing environment (sometimes referred to as a cloud environment, or a cloud) can be implemented in a variety of different ways to best suit different requirements. For example, in a public cloud environment, the underlying computing infrastructure is owned by an organization that makes its cloud services available to other organizations or to the general public. In contrast, a private cloud environment is generally intended solely for use by, or within, a single organization. A community cloud is intended to be shared by several organizations within a community; while a hybrid cloud comprise two or more types of cloud (e.g., private, community, or public) that are bound together by data and application portability.

(186) Generally, a cloud computing model enables some of those responsibilities which previously may have been provided by an organization's own information technology department, to instead be delivered as service layers within a cloud environment, for use by consumers (either within or external to the organization, according to the cloud's public/private nature). Depending on the particular implementation, the precise definition of components or features provided by or within each cloud service layer can vary, but common examples include: Software as a Service (SaaS), in which consumers use software applications that are running upon a cloud infrastructure, while a SaaS provider manages or controls the underlying cloud infrastructure and applications. Platform as a Service (PaaS), in which consumers can use software programming languages and development tools supported by a PaaS provider to develop, deploy, and otherwise control their own applications, while the PaaS provider manages or controls other aspects of the cloud environment (i.e., everything below the run-time execution environment). Infrastructure as a Service (IaaS), in which consumers can deploy and run arbitrary software applications, and/or

provision processing, storage, networks, and other fundamental computing resources, while an IaaS provider manages or controls the underlying physical cloud infrastructure (i.e., everything below the operating system layer). Database as a Service (DBaaS) in which consumers use a database server or Database Management System that is running upon a cloud infrastructure, while a DbaaS provider manages or controls the underlying cloud infrastructure and applications.

(187) The above-described basic computer hardware and software and cloud computing environment presented for purpose of illustrating the basic underlying computer components that may be employed for implementing the example embodiment(s). The example embodiment(s), however, are not necessarily limited to any particular computing environment or computing device configuration. Instead, the example embodiment(s) may be implemented in any type of system architecture or processing environment that one skilled in the art, in light of this disclosure, would understand as capable of supporting the features and functions of the example embodiment(s) presented herein.

(188) In the foregoing specification, embodiments of the invention have been described with reference to numerous specific details that may vary from implementation to implementation. The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense. The sole and exclusive indicator of the scope of the invention, and what is intended by the applicants to be the scope of the invention, is the literal and equivalent scope of the set of claims that issue from this application, in the specific form in which such claims issue, including any subsequent correction.

Claims

1. A method comprising: storing, within a computer memory, a first segmented array that comprises a first partial segment after a first one or more whole segments that contain a same particular amount of elements; storing, within the computer memory, a second segmented array that comprises a second partial segment after a second one or more whole segments that contain the same particular amount of elements; prepending, within the computer memory, the second one or more whole segments onto the first segmented array; and appending, within the computer memory, the second partial segment onto the first segmented array.
2. The method of claim 1 wherein the first segmented array and the second segmented array store respective subsets of values of a same particular column of a table that comprises multiple columns.
3. The method of claim 2 wherein the table contains a plurality of rows that represent elements of a graph selected from the group consisting of vertices of the graph and edges of the graph.
4. The method of claim 3 wherein: the graph is a distributed graph; the method further comprises receiving the second segmented array from a remote computer.
5. The method of claim 4 performed for rebalancing the distributed graph.
6. The method of claim 5 wherein: the method is performed by a particular computer; the rebalancing the distributed graph comprises moving a subset of vertices of the graph and a subset of edges of the graph from the remote computer to the particular computer.
7. The method of claim 4 wherein: the distributed graph is a property graph; each column of the multiple columns of the table stores values of a respective property of a plurality of properties of vertices or edges; the receiving the second segmented array from the remote computer comprises: receiving a value of a particular property that is not in the plurality of properties of vertices or edges, and generating a new column in the table that corresponds to the particular property.
8. The method of claim 2 wherein each column of the multiple columns of the table is stored in a respective segmented array.
9. The method of claim 1 further comprising a first execution thread accesses the first segmented array while a second execution thread simultaneously accesses the second segmented array.

10. The method of claim 1 further comprising configuring said same particular amount of elements to fit each whole segment of the first one or more whole segments and the second one or more whole segments into a respective cache line of a central processing unit (CPU).
 11. The method of claim 1 further comprising randomly accessing, based on a vector of distinct pointers to distinct segments in the first segmented array, a particular whole segment of the first one or more whole segments or the second one or more whole segments.
 12. The method of claim 1 further comprising generating two distinct vectors of distinct pointers to non-overlapping subsets of segments of the first segmented array.
 13. The method of claim 12 further comprising assigning each of the two distinct vectors of distinct pointers to respective distinct threads for parallel processing.
 14. The method of claim 1 wherein at least one selected from the group consisting of: the first one or more whole segments and the second one or more whole segments do not contain pointers, and the first segmented array and the second segmented array are not linked lists.
 15. The method of claim 1 wherein the first segmented array and the second segmented array can each contain at most one partial segment.
 16. The method of claim 15 wherein the at most one partial segment can only occur as a last segment of a segmented array.
 17. One or more computer-readable non-transitory media storing instructions that, when executed by one or more processors, cause: storing a first segmented array that comprises a first partial segment after a first one or more whole segments that contain a same particular amount of elements; storing a second segmented array that comprises a second partial segment after a second one or more whole segments that contain the same particular amount of elements; prepending the second one or more whole segments onto the first segmented array; and appending the second partial segment onto the first segmented array.
 18. The one or more computer-readable non-transitory media of claim 17 wherein the instructions further cause configuring said same particular amount of elements to fit each whole segment of the first one or more whole segments and the second one or more whole segments into a respective cache line of a central processing unit (CPU).
 19. The one or more computer-readable non-transitory media of claim 17 wherein the instructions further cause generating two distinct vectors of distinct pointers to non-overlapping subsets of the segments of first segmented array.
 20. The one or more computer-readable non-transitory media of claim 17 wherein at least one selected from the group consisting of: the first one or more whole segments and the second one or more whole segments do not contain pointers, and the first segmented array and the second segmented array are not linked lists.
-