

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411681
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Vishwanath; Manjunath et al.

Systems and methods for supporting rebootless firmware updates

Abstract

Embodiments of systems and methods to provide a firmware update to devices configured in a redundant configuration in an Information Handling System (IHS) are disclosed. In an illustrative, non-limiting embodiment, an IHS may include computer-executable instructions to receive a firmware update image associated with the firmware device, determine that the firmware update image possesses a critical sub-image portion, and schedule a firmware update to be performed in the future on the firmware device using the one firmware update image. The firmware update image conforms to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates.

Inventors: Vishwanath; Manjunath (Bangalore, IN), Bisa; Rama Rao (Bangalore, IN), Gavvala; Pavan Kumar (Bangalore, IN), Am; Manjunath (Bangalore, IN), Pradhan; Shantanu Kumar (Bangalore, IN), R; Chandrasekhar (Bangalore, IN)

Applicant: Dell Products, L.P. (Round Rock, TX)

Family ID: 90360458

Assignee: Dell Products, L.P. (Round Rock, TX)

Appl. No.: 17/935589

Filed: September 27, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240103849 A1	Mar. 28, 2024

Publication Classification

Int. Cl.: G06F8/656 (20180101)

U.S. Cl.:

CPC **G06F8/656** (20180201);

Field of Classification Search

CPC: G06F (8/656)

USPC: 717/170

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10437581	12/2018	Patil	N/A	G06F 8/71
10776286	12/2019	Kulchytsky	N/A	G06F 8/65
10996942	12/2020	Bisa	N/A	G06F 9/4812
2006/0174242	12/2005	Zhu	717/172	G06F 8/65
2011/0099544	12/2010	Haramiishi	717/168	G06F 8/654
2013/0145141	12/2012	Han	713/2	G06F 9/4401
2017/0147322	12/2016	Vopni	N/A	G06F 8/654
2020/0218545	12/2019	Ganesan	N/A	G06F 9/4411
2021/0124828	12/2020	Jarnik	N/A	G06F 21/54
2021/0141646	12/2020	Amin	N/A	G06F 9/4416
2021/0397436	12/2020	Le	N/A	G06F 9/4411

Primary Examiner: Wheaton; Bradford F

Attorney, Agent or Firm: Fogarty LLP

Background/Summary

BACKGROUND

(1) As the value and use of information continues to increase, individuals and businesses seek additional ways to process and store information. One option available to users is Information Handling Systems (IHSs). An IHS generally processes, compiles, stores, and/or communicates information or data for business, personal, or other purposes thereby allowing users to take advantage of the value of the information. Because technology and information handling needs and requirements vary between different users or applications, IHSs may also vary regarding what information is handled, how the information is handled, how much information is processed, stored, or communicated, and how quickly and efficiently the information may be processed, stored, or communicated. The variations in IHSs allow for IHSs to be general or configured for a specific user or specific use such as financial transaction processing, airline reservations, enterprise data storage, or global communications. In addition, IHSs may include a variety of hardware and software components that may be configured to process, store, and communicate information and may include one or more computer systems, data storage systems, and networking systems.

(2) Various hardware components of an IHS may operate using firmware instructions. From time to time, it is expected that firmware utilized by hardware components of an IHS may be updated.

Such firmware updates may be made in order to modify the capabilities of a particular hardware component, such as to address security vulnerabilities or to adapt the operations of the hardware component to a specific computing task. When firmware updates are made to a hardware component of an IHS, it is preferable that the IHS experience no downtime and with minimal degradation in the performance of the IHS.

(3) Nowadays, software updates are typically made available on one or more download sites as soon as the software provider can produce them. In this manner, software providers can be more responsive to critical flaws, security concerns, and general customer needs. To update software, a customer would query an update site for software updates, and download and install the software update if available. For example, a typical network-based software update procedure may include the steps of issuing a request over a network to a software provider's download site (e.g., update source) for a software update applicable to the client computer. The update source responds to the client computer with the software update requested by the client computer in the update request. After the client computer has received the software update, the client computer installs the received software update.

(4) One benefit of updating software in such a manner is the reduced cost associated with producing and distributing software updates. Additionally, software updates can now be performed more frequently, especially those that address critical issues and security. Still further, a computer user has greater control as to when and which software updates should be installed on the client computer.

SUMMARY

(5) Embodiments of systems and methods to provide a firmware update to devices configured in a redundant configuration in an Information Handling System (IHS) are disclosed. In an illustrative, non-limiting embodiment, an IHS may include computer-executable instructions to receive a firmware update image associated with the firmware device, determine that the firmware update image possesses a critical sub-image portion, and schedule a firmware update to be performed in the future on the firmware device using the one firmware update image. The firmware update image conforms to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates.

(6) According to another embodiment, a rebootless firmware update method includes the steps of receiving a firmware update image associated with the firmware device, the firmware update image conforming to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates, determining that the firmware update image possesses a critical sub-image portion, and scheduling a firmware update to be performed in the future on the firmware device using the one firmware update image.

(7) According to yet another embodiment, a memory storage device is configured with program instructions that, upon execution by an Information Handling System (IHS), cause the IHS to receive a firmware update image associated with the firmware device, the firmware update image conforming to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates, determine that the firmware update image possesses a critical sub-image portion, and schedule a firmware update to be performed in the future on the firmware device using the one firmware update image.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present invention(s) is/are illustrated by way of example and is/are not limited by the accompanying figures. Elements in the figures are illustrated for simplicity and clarity, and have not necessarily been drawn to scale.

- (2) FIGS. 1A and 1B illustrate certain components of a chassis comprising one or more compute sleds and one or more storage sleds that may be configured to implement the systems and methods described according to one embodiment of the present disclosure.
- (3) FIG. 2 illustrates an example of an IHS configured to implement systems and methods described herein according to one embodiment of the present disclosure.
- (4) FIG. 3 is a diagram illustrating several components of an IHS showing how those components may communicate with one another for implementing an example rebootless firmware update system according to one embodiment of the present disclosure.
- (5) FIG. 4 illustrates an example the various fields of an OEM command that may be used to query the firmware device for version information for each of its installed sub-components.
- (6) FIG. 5 illustrates an example sub-component version acquisition method that may be used to acquire versioning information for the sub-components of installed firmware in the firmware devices configured in an IHS according to one embodiment of the present disclosure.
- (7) FIG. 6 illustrates an example rebootless firmware update method depicting how a firmware device configured in an IHS may be updated according to one embodiment of the present disclosure.

DETAILED DESCRIPTION

- (8) The present disclosure is described with reference to the attached figures. The figures are not drawn to scale, and they are provided merely to illustrate the disclosure. Several aspects of the disclosure are described below with reference to example applications for illustration. It should be understood that numerous specific details, relationships, and methods are set forth to provide an understanding of the disclosure. The present disclosure is not limited by the illustrated ordering of acts or events, as some acts may occur in different orders and/or concurrently with other acts or events. Furthermore, not all illustrated acts or events are required to implement a methodology in accordance with the present disclosure.
- (9) Firmware updates of server components is an important aspect of the life cycle management of an Information Handling System (IHS), (e.g., server, host, etc.). Traditional means of updating IHS devices have involved migrating the workloads running on the host Operating System (OS), creating a reboot job, rebooting the IHS, and performing the firmware update. Additionally, the IHS is again rebooted to activate the new firmware on the IHS components. This process, however, may not be customer friendly as the IHS is often required to be down for the firmware update process, thus impacting business. Because IHSs are forced to reboot to perform the firmware updates, customers often wait for its maintenance cycle (e.g., maintenance window) to update the IHS components, thus missing the new firmware features, security fixes, performance improvements, and the like. As such, rebootless updates may be an important aspect of efficient computer operations. Using rebootless updates, users may be enabled with performing the updates without rebooting the servers and get more useful features above what today's industry specifications can provide.
- (10) Customers often upgrade the firmware in the IHSs of a data center for assorted reasons, such as to meet compliance policies, to take advantage of new features, enhancements to the IHS, deploy security fixes, and the like. IHSs that are NVMe-MI/PLDM Specification compliant can take advantage of updating firmware to all IHSs in a system or in a cluster without rebooting the IHSs. Devices that support Platform Level Data Model (PLDM) offer an option for a Remote Access Controller (RAC) to update the firmware without rebooting the IHS. Thus, downtime is often not incurred during the firmware update process. The RAC may be configured to provide out-of-band management facilities for an IHS, even if it is powered off, or powered down to a standby state. The RAC may include a processor, memory, and an out-of-band network interface separate from and physically isolated from an in-band network interface of the IHS, and/or other embedded resources. In certain embodiments, the RAC may include or may be part of a Remote Access Controller (e.g., a DELL Remote Access Controller (DRAC) or an Integrated DRAC (iDRAC)).

(11) The RAC may support rebootless firmware updates for devices, such as non-volatile storage (e.g., hard disks, Solid State Drives (SSDs), etc.), Network Interface Cards (NICs), Graphical Processing Units (GPUs), RACs, Hardware RAID (HWRAID) devices, and the like. With the rebootless feature, when a firmware update image is uploaded using a RAC user interface, all the devices supported by the firmware update image may be automatically selected and updated using rebootless update methods in the real-time without rebooting the IHS. The RAC may implement a Platform Management Components Intercommunication (PMCI) interface stack that is provided by the Distributed Management Task Force (DMTF), and specifies a Management Component Transport Protocol (MCTP) specifying how data travels over certain physical layers, such as the peripheral component interconnect express (PCIe) and I2C/SMBus. Additionally, the PMCI interface stack may further include the Platform Level Data Model (PLDM) protocol that enables information to travel over the MCTP transport layer and can be used for platform management, such as firmware updates.

(12) As mentioned previously, the NVMe-MI and PLDM firmware update specifications provide a way for the RAC to transfer the firmware update images and activate the new firmware update images without the need for a power cycle/server reboot. Though transfer of the firmware update image is often possible without the need for a reboot, it may not necessarily be able to be activated without being rebooted. The firmware device which claims rebootless update capability before the firmware image transfer, can change its capability claims after the firmware update image is completely transferred and verified. Several reasons for such a scenario may include the fact that one or more critical component image portions might have changed in the new firmware update image relative to the existing firmware image on the device and as such, the firmware devices cannot activate the new firmware update image without a power cycle/reboot of the server.

(13) Conventionally, when a firmware update image is uploaded to the RAC, it reads the rebootless/self-contained image activation capability of the firmware device, and schedules a firmware update to be performed immediately if the device claims a rebootless update capability. The user may think they can update the firmware without server reboot given its rebootless update capability claim. But sometimes after the image is transferred, the firmware device may request for reboot of the server, thus potentially mis-informing the user unintentionally, and leading to a bad customer experience.

(14) This problem may be exacerbated by the fact that currently available firmware update images may be up to 40 Mega-bytes in size, and when transferred over an I2C communication link, could consume a relatively large amount of time given that I2C communication links often require approximately 53 minutes to transfer a 100 Kilo-byte image. Other firmware devices using the same I2C communication link could potentially be starved for tens of minutes. Such bottleneck-type behavior could result in causing the fans to spin full speed, and stopping some, most, or all the monitoring and configuration activities during the firmware update. As such, if the firmware device changes its claim on rebootless capability, it could lead to a bad customer experience. With the current industry specifications, the customers cannot be informed when a particular firmware update will need a reboot with a high degree of consistency. Otherwise, such scenarios could lead to bad customer experience and defeat the zero downtime criteria for high availability (HA) based workloads. As will be described in detail herein below, embodiments the present disclosure provide a rebootless firmware update system and method that receives a firmware update image associated with the firmware device, determines whether the firmware update image possesses a critical component update portion, and if so, schedules a firmware update to be performed in the future on the firmware device using the one firmware update image.

(15) FIGS. 1A and 1B illustrate certain components of a chassis **100** comprising one or more compute sleds **105a-n** and one or more storage sleds **115a-n** that may be configured to implement the systems and methods described according to one embodiment of the present disclosure. Embodiments of chassis **100** may include a wide variety of hardware configurations in which one

or more sleds **105a-n**, **115a-n** are installed in chassis **100**. Such variations in hardware configuration may result from chassis **100** being factory assembled to include components specified by a customer that has contracted for manufacture and delivery of chassis **100**. Upon delivery and deployment of a chassis **100**, the chassis **100** may be modified by replacing and/or adding various hardware components, in addition to replacement of the removable sleds **105a-n**, **115a-n** that are installed in the chassis. In addition, once the chassis **100** has been deployed, firmware used by individual hardware components of the sleds **105a-n**, **115a-n**, or by other hardware components of chassis **100**, may be modified in order to update the operations that are supported by these hardware components.

(16) Chassis **100** may include one or more bays that each receive an individual sled (that may be additionally or alternatively referred to as a tray, blade, and/or node), such as compute sleds **105a-n** and storage sleds **115a-n**. Chassis **100** may support a variety of different numbers (e.g., 4, 8, 16, 32), sizes (e.g., single-width, double-width) and physical configurations of bays. Embodiments may include additional types of sleds that provide various storage, power and/or processing capabilities. For instance, sleds installable in chassis **100** may be dedicated to providing power management or networking functions. Sleds may be individually installed and removed from the chassis **100**, thus allowing the computing and storage capabilities of a chassis to be reconfigured by swapping the sleds with diverse types of sleds, in some cases at runtime without disrupting the ongoing operations of the other sleds installed in the chassis **100**.

(17) Multiple chassis **100** may be housed within a rack. Data centers may utilize large numbers of racks, with various different types of chassis installed in various configurations of racks. The modular architecture provided by the sleds, chassis and racks allow for certain resources, such as cooling, power and network bandwidth, to be shared by the compute sleds **105a-n** and storage sleds **115a-n**, thus providing efficiency improvements and supporting greater computational loads. For instance, certain computational tasks, such as computations used in machine learning and other artificial intelligence systems, may utilize computational and/or storage resources that are shared within an IHS, within an individual chassis **100** and/or within a set of IHSs that may be spread across multiple chassis of a data center.

(18) Implementing computing systems that span multiple processing components of chassis **100** is aided by high-speed data links between these processing components, such as PCIe connections that form one or more distinct PCIe switch fabrics that are implemented by PCIe switches **135a-n**, **165a-n** installed in the sleds **105a-n**, **115a-n** of the chassis. These high-speed data links may be used to support algorithm implementations that span multiple processing, networking, and storage components of an IHS and/or chassis **100**. For instance, computational tasks may be delegated to a specific processing component of an IHS, such as to a hardware accelerator **185a-n** that may include one or more programmable processors that operate separate from the main CPUs **170a-n** of computing sleds **105a-n**. In various embodiments, such hardware accelerators **185a-n** may include DPUs (Data Processing Units), GPUs (Graphics Processing Units), SmartNICs (Smart Network Interface Card) and/or FPGAs (Field Programmable Gate Arrays). These hardware accelerators **185a-n** operate according to firmware instructions that may be occasionally updated, such as to adapt the capabilities of the respective hardware accelerators **185a-n** to specific computing tasks.

(19) Chassis **100** may be installed within a rack structure that provides at least a portion of the cooling utilized by the sleds **105a-n**, **115a-n** installed in chassis **100**. In supporting airflow cooling, a rack may include one or more banks of cooling fans **130** that may be operated to ventilate heated air from within the chassis **100** that is housed within the rack. The chassis **100** may alternatively or additionally include one or more cooling fans **130** that may be similarly operated to ventilate heated air away from sleds **105a-n**, **115a-n** installed within the chassis. In this manner, a rack and a chassis **100** installed within the rack may utilize various configurations and combinations of cooling fans **130** to cool the sleds **105a-n**, **115a-n** and other components housed within chassis **100**.

(20) The sleds **105a-n**, **115a-n** may be individually coupled to chassis **100** via connectors that

correspond to the bays provided by the chassis **100** and that physically and electrically couple an individual sled to a backplane **160**. Chassis backplane **160** may be a printed circuit board that includes electrical traces and connectors that are configured to route signals between the various components of chassis **100** that are connected to the backplane **160** and between different components mounted on the printed circuit board of the backplane **160**. In the illustrated embodiment, the connectors for use in coupling sleds **105a-n**, **115a-n** to backplane **160** include PCIe couplings that support high-speed data links with the sleds **105a-n**, **115a-n**. In various embodiments, backplane **160** may support diverse types of connections, such as cables, wires, midplanes, connectors, expansion slots, and multiplexers. In certain embodiments, backplane **160** may be a motherboard that includes various electronic components installed thereon. Such components installed on a motherboard backplane **160** may include components that implement all or part of the functions described with regard to the SAS (Serial Attached SCSI) expander **150**, I/O controllers **145**, network controller **140**, chassis management controller **125** and power supply unit **135**.

(21) In certain embodiments, each individual sled **105a-n**, **115a-n** may be an IHS such as described with regard to IHS **200** of FIG. 2. Sleds **105a-n**, **115a-n** may individually or collectively provide computational processing resources that may be used to support a variety of e-commerce, multimedia, business, and scientific computing applications, such as artificial intelligence systems provided via cloud computing implementations. Sleds **105a-n**, **115a-n** are typically configured with hardware and software that provide leading-edge computational capabilities. Accordingly, services that are provided using such computing capabilities are typically provided as high-availability systems that operate with minimum downtime.

(22) In high-availability computing systems, such as may be implemented using embodiments of chassis **100**, any downtime that can be avoided is preferred. As described above, firmware updates are expected in the administration and operation of data centers, but it is preferable to avoid any downtime in making such firmware updates. For instance, in updating the firmware of the individual hardware components of the chassis **100**, it is preferable that such updates can be made without having to reboot the chassis. As described in additional detail below, it is also preferable that updates to the firmware of individual hardware components of sleds **105a-n**, **115a-n** be likewise made without having to reboot the respective sled of the hardware component that is being updated.

(23) As illustrated, each sled **105a-n**, **115a-n** includes a respective remote access controller (RAC) **110a-n**, **120a-n**. As described in additional detail with regard to FIG. 2, remote access controller **110a-n**, **120a-n** provides capabilities for remote monitoring and management of a respective sled **105a-n**, **115a-n** and/or of chassis **100**. In support of these monitoring and management functions, remote access controllers **110a-n** may utilize both in-band and sideband (i.e., out-of-band) communications with various managed components of a respective sled **105a-n**, **115a-n** and chassis **100**. Remote access controllers **110a-n**, **120a-n** may collect diverse types of sensor data, such as collecting temperature sensor readings that are used in support of airflow cooling of the chassis **100** and the sled **105a-n**, **115a-n**. In addition, each remote access controller **110a-n**, **120a-n** may implement various monitoring and administrative functions related to a respective sled **105a-n**, **115a-n**, where these functions may be implemented using sideband bus connections with various internal components of the chassis **100** and of the respective sleds **105a-n**, **115a-n**. As described in additional detail below, in various embodiments, these capabilities of the remote access controllers **110a-n**, **120a-n** may be utilized in updating the firmware of hardware components of chassis **100** and/or of hardware components of the sleds **105a-n**, **115a-n**, without having to reboot the chassis or any of the sleds **105a-n**, **115a-n**.

(24) The remote access controllers **110a-n**, **120a-n** that are present in chassis **100** may support secure connections with a remote management interface **101**. In some embodiments, remote management interface **101** provides a remote administrator with various capabilities for remotely

administering the operation of an IHS, including initiating updates to the firmware used by hardware components installed in the chassis **100**. For example, remote management interface **101** may provide capabilities by which an administrator can initiate updates to all of the storage drives **175a-n** installed in a chassis **100**, or to all of the storage drives **175a-n** of a particular model or manufacturer. In some instances, remote management interface **101** may include an inventory of the hardware, software, and firmware of chassis **100** that is being remotely managed through the operation of the remote access controllers **110a-n**, **120a-n**. The remote management interface **101** may also include various monitoring interfaces for evaluating telemetry data collected by the remote access controllers **110a-n**, **120a-n**. In some embodiments, remote management interface **101** may communicate with remote access controllers **110a-n**, **120a-n** via a protocol such as the Redfish remote management interface.

(25) In the illustrated embodiment, chassis **100** includes one or more compute sleds **105a-n** that are coupled to the backplane **160** and installed within one or more bays or slots of chassis **100**. Each of the individual compute sleds **105a-n** may be an IHS, such as described with regard to FIG. 2. Each of the individual compute sleds **105a-n** may include various different numbers and types of processors that may be adapted to performing specific computing tasks. In the illustrated embodiment, each of the compute sleds **105a-n** includes a PCIe switch **135a-n** that provides access to a hardware accelerator **185a-n**, such as the described DPUs, GPUs, Smart NICs and FPGAs, which may be programmed and adapted for specific computing tasks, such as to support machine learning or other artificial intelligence systems. As described in additional detail below, compute sleds **105a-n** may include a variety of hardware components, such as hardware accelerator **185a-n** and PCIe switches **135a-n**, that operate using firmware that may be occasionally updated.

(26) As illustrated, chassis **100** includes one or more storage sleds **115a-n** that are coupled to the backplane **160** and installed within one or more bays of chassis **100** in a similar manner to compute sleds **105a-n**. Each of the individual storage sleds **115a-n** may include various different numbers and types of storage devices. As described in additional detail with regard to FIG. 2, a storage sled **115a-n** may be an IHS **200** that includes multiple solid-state drives (SSDs) **175a-n**, where the individual storage drives **175a-n** may be accessed through a PCIe switch **165a-n** of the respective storage sled **115a-n**.

(27) As illustrated, a storage sled **115a** may include one or more DPUs (Data Processing Units) **190** that provide access to and manage the operations of the storage drives **175a** of the storage sled **115a**. Use of a DPU **190** in this manner provides low-latency and high-bandwidth access to numerous SSDs **175a**. These SSDs **175a** may be utilized in parallel through NVMe transmissions that are supported by the PCIe switch **165a** that connects the SSDs **175a** to the DPU **190**. In some instances, PCIe switch **165a** may be an integrated component of a DPU **190**. The immense data storage and retrieval capabilities provided by such storage sled **115a** implementations may be harnessed by offloading storage operations directed at storage drives **175a** to a DPU **190**, and thus without relying on the main CPU of the storage sled, or of any other component of chassis **100**. As indicated in FIG. 1, chassis **100** may also include one or more storage sleds **115n** that provide access to storage drives **175n** via a storage controller **195**. In some embodiments, storage controller **195** may provide support for RAID (Redundant Array of Independent Disks) configurations of logical and physical storage drives, such as storage drives provided by storage sled **115n**. In some embodiments, storage controller **195** may be a HBA (Host Bus Adapter) that provides more limited capabilities in accessing storage drives **175n**.

(28) In addition to the data storage capabilities provided by storage sleds **115a-n**, chassis **100** may provide access to other storage resources that may be installed in chassis **100** and/or may be installed elsewhere within a rack that houses the chassis **100**. In certain scenarios, such storage resources (e.g., JBOD **155**) may be accessed via a SAS expander **150** that is coupled to the backplane **160** of the chassis **100**. The SAS expander **150** may support connections to a number of JBOD (Just a Bunch of Disks) storage resources **155** that, in some instances, may be configured

and managed individually and without implementing data redundancy across the various drives. The additional storage resources may also be at various other locations within a datacenter in which chassis **100** is installed.

(29) In light of the various manners in which storage drives **175a-n**, **155** may be coupled to chassis **100**, a wide variety of different storage topologies may be supported. Through these supported topologies, storage drives **175a-n**, **155** may be logically organized into clusters or other groupings that may be collectively tasked and managed. In some instances, a chassis **100** may include numerous storage drives **175a-n**, **155** that are identical, or nearly identical, such as arrays of SSDs of the same manufacturer and model. Accordingly, any firmware updates to storage drives **175a-n**, **155** requires the updates to be applied within each of these topologies being supported by the chassis **100**. Despite the large number of different storage drive topologies that may be supported by an individual chassis **100**, the firmware used by each of these storage devices **175a-n**, **155** may be occasionally updated. In some instances, firmware updates may be limited to a single storage drive, but in other instances, firmware updates may be initiated for a large number of storage drives, such as for all SSDs installed in chassis **100**.

(30) As illustrated, the chassis **100** of FIG. **1** includes a network controller **140** that provides network access to the sleds **105a-n**, **115a-n** installed within the chassis. Network controller **140** may include various switches, adapters, controllers, and couplings used to connect chassis **100** to a network, either directly or via additional networking components and connections provided via a rack in which chassis **100** is installed. Network controller **140** operates according to firmware instructions that may be occasionally updated.

(31) Chassis **100** may similarly include a power supply unit **135** that provides the components of the chassis with various levels of DC power from an AC power source or from power delivered via a power system provided by a rack within which chassis **100** may be installed. In certain embodiments, power supply unit **135** may be implemented within a sled that may provide chassis **100** with redundant, hot-swappable power supply units. Power supply unit **135** may operate according to firmware instructions that may be occasionally updated.

(32) Chassis **100** may also include various I/O controllers **145** that may support various I/O ports, such as USB ports that may be used to support keyboard and mouse inputs and/or video display capabilities. Each of the I/O controllers **145** may operate according to firmware instructions that may be occasionally updated. Such I/O controllers **145** may be utilized by the chassis management controller **125** to support various KVM (Keyboard, Video and Mouse) **125a** capabilities that provide administrators with the ability to interface with the chassis **100**. The chassis management controller **125** may also include a storage module **125c** that provides capabilities for managing and configuring certain aspects of the storage devices of chassis **100**, such as the storage devices provided within storage sleds **115a-n** and within the JBOD **155**.

(33) In addition to providing support for KVM **125a** capabilities for administering chassis **100**, chassis management controller **125** may support various additional functions for sharing the infrastructure resources of chassis **100**. In some scenarios, chassis management controller **125** may implement tools for managing the power supply unit **135**, network controller **140** and airflow cooling fans **130** that are available via the chassis **100**. As described, the airflow cooling fans **130** utilized by chassis **100** may include an airflow cooling system that is provided by a rack in which the chassis **100** may be installed and managed by a cooling module **125b** of the chassis management controller **125**.

(34) For purposes of this disclosure, an IHS may include any instrumentality or aggregate of instrumentalities operable to compute, calculate, determine, classify, process, transmit, receive, retrieve, originate, switch, store, display, communicate, manifest, detect, record, reproduce, handle, or utilize any form of information, intelligence, or data for business, scientific, control, or other purposes. For example, an IHS may be a personal computer (e.g., desktop or laptop), tablet computer, mobile device (e.g., Personal Digital Assistant (PDA) or smart phone), server (e.g., blade

server or rack server), a network storage device, or any other suitable device and may vary in size, shape, performance, functionality, and price. An IHS may include Random Access Memory (RAM), one or more processing resources such as a Central Processing Unit (CPU) or hardware or software control logic, Read-Only Memory (ROM), and/or other types of nonvolatile memory. Additional components of an IHS may include one or more disk drives, one or more network ports for communicating with external devices as well as various I/O devices, such as a keyboard, a mouse, touchscreen, and/or a video display. As described, an IHS may also include one or more buses operable to transmit communications between the various hardware components. An example of an IHS is described in more detail below.

(35) FIG. 2 illustrates an example of an IHS **200** configured to implement systems and methods described herein according to one embodiment of the present disclosure. It should be appreciated that although the embodiments described herein may describe an IHS that is a compute sled or similar computing component that may be deployed within the bays of a chassis, a variety of other types of IHSs, such as laptops and portable devices, may also operate according to embodiments described herein. In the illustrative embodiment of FIG. 2, IHS **200** may be a computing component, such as sled **105a-n** or other type of server, such as a 1RU server installed within a 2RU chassis, which is configured to share infrastructure resources provided within a chassis **100**.

(36) IHS **200** may utilize one or more system processors **205**, that may be referred to as CPUs (central processing units). In some embodiments, CPUs **205** may each include a plurality of processing cores that may be separately delegated with computing tasks. Each of the CPUs **205** may be individually designated as a main processor and as a co-processor, where such designations may be based on delegation of specific types of computational tasks to a CPU **205**. In some embodiments, CPUs **205** may each include an integrated memory controller that may be implemented directly within the circuitry of each CPU **205**. In some embodiments, a memory controller may be a separate integrated circuit that is located on the same die as the CPU **205**. Each memory controller may be configured to manage the transfer of data to and from a system memory **210** of the IHS, in some cases using a high-speed memory bus **205a**. The system memory **210** is coupled to CPUs **205** via one or more memory buses **205a** that provide the CPUs **205** with high-speed memory used in the execution of computer program instructions by the CPUs **205**.

Accordingly, system memory **210** may include memory components, such as static RAM (SRAM), dynamic RAM (DRAM), NAND Flash memory, suitable for supporting high-speed memory operations by the CPUs **205**. In certain embodiments, system memory **210** may combine persistent non-volatile memory and volatile memory.

(37) In certain embodiments, the system memory **210** may be comprised of multiple removable memory modules. The system memory **210** of the illustrated embodiment includes removable memory modules **210a-n**. Each of the removable memory modules **210a-n** may correspond to a printed circuit board memory socket that receives a removable memory module **210a-n**, such as a DIMM (Dual In-line Memory Module), that can be coupled to the socket and then decoupled from the socket as needed, such as to upgrade memory capabilities or to replace faulty memory modules. Other embodiments of IHS system memory **210** may be configured with memory socket interfaces that correspond to diverse types of removable memory module form factors, such as a Dual In-line Package (DIP) memory, a Single In-line Pin Package (SIPP) memory, a Single In-line Memory Module (SIMM), and/or a Ball Grid Array (BGA) memory.

(38) IHS **200** may utilize a chipset that may be implemented by integrated circuits that are connected to each CPU **205**. All or portions of the chipset may be implemented directly within the integrated circuitry of an individual CPU **205**. The chipset may provide the CPU **205** with access to a variety of resources accessible via one or more in-band buses. IHS **200** may also include one or more I/O ports **215** that may be used to couple the IHS **200** directly to other IHSs, storage resources, diagnostic tools, and/or other peripheral components. A variety of additional components may be coupled to CPUs **205** via a variety of in-line buses. For instance, CPUs **205** may also be

coupled to a power management unit **220** that may interface with a power system of the chassis **100** in which IHS **200** may be installed. In addition, CPUs **205** may collect information from one or more sensors **225** via a management bus.

(39) In certain embodiments, IHS **200** may operate using a BIOS (Basic Input/Output System) that may be stored in a non-volatile memory accessible by the CPUs **205**. The BIOS may provide an abstraction layer by which the operating system of the IHS **200** interfaces with hardware components of the IHS. Upon powering or restarting IHS **200**, CPUs **205** may utilize BIOS instructions to initialize and test hardware components coupled to the IHS, including both components permanently installed as components of the motherboard of IHS **200** and removable components installed within various expansion slots supported by the IHS **200**. The BIOS instructions may also load an operating system for execution by CPUs **205**. In certain embodiments, IHS **200** may utilize Unified Extensible Firmware Interface (UEFI) in addition to or instead of a BIOS. In certain embodiments, the functions provided by a BIOS may be implemented, in full or in part, by the remote access controller **230**.

(40) In some embodiments, IHS **200** may include a TPM (Trusted Platform Module) that may include various registers, such as platform configuration registers, and a secure storage, such as an NVRAM (Non-Volatile Random-Access Memory). The TPM may also include a cryptographic processor that supports various cryptographic capabilities. In IHS embodiments that include a TPM, a pre-boot process implemented by the TPM may utilize its cryptographic capabilities to calculate hash values that are based on software and/or firmware instructions utilized by certain core components of IHS, such as the BIOS and boot loader of IHS **200**. These calculated hash values may then be compared against reference hash values that were previously stored in a secure non-volatile memory of the IHS, such as during factory provisioning of IHS **200**. In this manner, a TPM may establish a root of trust that includes core components of IHS **200** that are validated as operating using instructions that originate from a trusted source.

(41) As illustrated, CPUs **205** may be coupled to a network controller **240**, such as provided by a Network Interface Controller (NIC) card that provides IHS **200** with communications via one or more external networks, such as the Internet, a LAN, or a WAN. In some embodiments, network controller **240** may be a replaceable expansion card or adapter that is coupled to a connector (e.g., PCIe connector of a motherboard, backplane, midplane, etc.) of IHS **200**. In some embodiments, network controller **240** may support high-bandwidth network operations by the IHS **200** through a PCIe interface that is supported by the chipset of CPUs **205**. Network controller **240** may operate according to firmware instructions that may be occasionally updated.

(42) As indicated in FIG. 2, in some embodiments, CPUs **205** may be coupled to a PCIe card **255** that includes two PCIe switches **265a-b** that operate as I/O controllers for PCIe communications, such as TLPs (Transaction Layer Packets), that are transmitted between the CPUs **205** and PCIe devices and systems coupled to IHS **200**. Whereas the illustrated embodiment of FIG. 2 includes two CPUs **205** and two PCIe switches **265a-b**, different embodiments may operate using different numbers of CPUs and PCIe switches. In addition to serving as I/O controllers that route PCIe traffic, PCIe switches **265a-b** include switching logic that can be used to expand the number of PCIe connections that are supported by CPUs **205**. PCIe switches **265a-b** may multiply the number of PCIe lanes available to CPUs **205**, thus allowing more PCIe devices to be connected to CPUs **205**, and for the available PCIe bandwidth to be allocated with greater granularity. Each of the PCIe switches **265a-b** may operate according to firmware instructions that may be occasionally updated.

(43) Using the available PCIe lanes, the PCIe switches **265a-b** may be used to implement a PCIe switch fabric. Also through this switch fabric, PCIe NVMe (Non-Volatile Memory Express) transmission may be supported and utilized in high-speed communications with SSDs, such as storage drives **235a-b**, of the IHS **200**. Also through this switch fabric, PCIe VDM (Vendor Defined Messaging) may be supported and utilized in managing PCIe-compliant hardware components of the IHS **200**, such as in updating the firmware utilized by the hardware components.

(44) As indicated in FIG. 2, IHS **200** may support storage drives **235a-b** in various topologies, in the same manner as described with regard to the chassis **100** of FIG. 1. In the illustrated embodiment, storage drives **235a** are accessed via a hardware accelerator **250**, while storage drives **235b** are accessed directly via PCIe switch **265b**. In some embodiments, the storage drives **235a-b** of IHS **200** may include a combination of both SSD and magnetic disk storage drives. In other embodiments, all of the storage drives **235a-b** of IHS **200** may be identical, or nearly identical. In all embodiments, storage drives **235a-b** operate according to firmware instructions that may be occasionally updated.

(45) As illustrated, PCIe switch **265a** is coupled via a PCIe link to a hardware accelerator **250**, such as a DPU, SmartNIC, GPU and/or FPGA, that may be connected to the IHS via a removable card or baseboard that couples to a PCIe connector of the IHS **200**. In some embodiments, hardware accelerator **250** includes a programmable processor that can be configured for offloading functions from CPUs **205**. In some embodiments, hardware accelerator **250** may include a plurality of programmable processing cores and/or hardware accelerators, which may be used to implement functions used to support devices coupled to the IHS **200**. In some embodiments, the processing cores of hardware accelerator **250** include ARM (advanced RISC (reduced instruction set computing) machine) processing cores. In other embodiments, the cores of the DPUs may include MIPS (microprocessor without interlocked pipeline stages) cores, RISC-V cores, or CISC (complex instruction set computing) (i.e., x86) cores. Hardware accelerator may operate according to firmware instructions that may be occasionally updated.

(46) In the illustrated embodiment, the programmable capabilities of hardware accelerator **250** implement functions used to support storage drives **235a**, such as SSDs. In such storage drive topologies, hardware accelerator **250** may implement processing of PCIe NVMe communications with SSDs **235a**, thus supporting high-bandwidth connections with these SSDs. Hardware accelerator **250** may also include one or more memory devices used to store program instructions executed by the processing cores and/or used to support the operation of SSDs **235a** such as in implementing cache memories and buffers utilized to support high-speed operation of these storage drives, and in some cases may be used to provide high-availability and high-throughput implementations of the read, write and other I/O operations that are supported by these storage drives **235a**. In other embodiments, hardware accelerator **250** may implement operations in support of other types of devices and may similarly support high-bandwidth PCIe connections with these devices. For instance, in various embodiments, hardware accelerator **250** may support high-bandwidth connections, such as PCIe connections, with networking devices in implementing functions of a network switch, compression and codec functions, virtualization operations or cryptographic functions.

(47) As illustrated in FIG. 2, PCIe switches **265a-b** may also support PCIe couplings with one or more GPUs (Graphics Processing Units) **260**. Embodiments may include one or more GPU cards, where each GPU card is coupled to one or more of the PCIe switches **265a-b**, and where each GPU card may include one or more GPUs **260**. In some embodiments, PCIe switches **265a-b** may transfer instructions and data for generating video images by the GPUs **260** to and from CPUs **205**. Accordingly, GPUs **260** may include one or more hardware-accelerated processing cores that are optimized for performing streaming calculation of vector data, matrix data and/or other graphics data, thus supporting the rendering of graphics for display on devices coupled either directly or indirectly to IHS **200**. In some instances, GPUs may be utilized as programmable computing resources for offloading other functions from CPUs **205**, in the same manner as hardware accelerator **250**. GPUs **260** may operate according to firmware instructions that may be occasionally updated.

(48) As illustrated in FIG. 2, PCIe switches **265a-b** may support PCIe connections in addition to those utilized by GPUs **260** and hardware accelerator **250**, where these connections may include PCIe links of one or more lanes. For instance, PCIe connectors **245** supported by a printed circuit

board of IHS **200** may allow various other systems and devices to be coupled to IHS. Through couplings to PCIe connectors **245**, a variety of data storage devices, graphics processors and network interface cards may be coupled to IHS **200**, thus supporting a wide variety of topologies of devices that may be coupled to the IHS **200**.

(49) As described, IHS **200** includes a remote access controller **230** that supports remote management of IHS **200** and of various internal components of IHS **200**. In certain embodiments, remote access controller **230** may operate from a different power plane from the processors **205** and other components of IHS **200**, thus allowing the remote access controller **230** to operate, and manage tasks to proceed, while the processing cores of IHS **200** are powered off. Various functions provided by the BIOS, including launching the operating system of the IHS **200**, and/or functions of a TPM may be implemented or supplemented by the remote access controller **230**. In some embodiments, the remote access controller **230** may perform various functions to verify the integrity of the IHS **200** and its hardware components prior to initialization of the operating system of IHS **200** (i.e., in a bare-metal state). In some embodiments, certain operations of the remote access controller **230**, such as the operations described herein for updating firmware used by managed hardware components of IHS **200**, may operate using validated instructions, and thus within the root of trust of IHS **200**.

(50) In some embodiments, remote access controller **230** may include a service processor **230a**, or specialized microcontroller, which operates management software that supports remote monitoring and administration of IHS **200**. The management operations supported by remote access controller **230** may be remotely initiated, updated, and monitored via a remote management interface **101**, such as described with regard to FIG. **1**. Remote access controller **230** may be installed on the motherboard of IHS **200** or may be coupled to IHS **200** via an expansion slot or other connector provided by the motherboard. In some instances, the management functions of the remote access controller **230** may utilize information collected by various managed sensors **225** located within the IHS. For instance, temperature data collected by sensors **225** may be utilized by the remote access controller **230** in support of closed-loop airflow cooling of the IHS **200**. As indicated, remote access controller **230** may include a secured memory **230e** for exclusive use by the remote access controller in support of management operations.

(51) In some embodiments, remote access controller **230** may implement monitoring and management operations using MCTP (Management Component Transport Protocol) messages that may be communicated to managed devices **205**, **235a-b**, **240**, **250**, **255**, **260** via management connections supported by a sideband bus **253**. In some embodiments, the remote access controller **230** may additionally or alternatively use MCTP messaging to transmit Vendor Defined Messages (VDMs) via the in-line PCIe switch fabric supported by PCIe switches **265a-b**. In some instances, the sideband management connections supported by remote access controller **230** may include PLDM (Platform Level Data Model) management communications with the managed devices **205**, **235a-b**, **240**, **250**, **255**, **260** of IHS **200**.

(52) As illustrated, remote access controller **230** may include a network adapter **230c** that provides the remote access controller with network access that is separate from the network controller **240** utilized by other hardware components of the IHS **200**. Through secure connections supported by network adapter **230c**, remote access controller **230** communicates management information with remote management interface **101**. In support of remote monitoring functions, network adapter **230c** may support connections between remote access controller **230** and external management tools using wired and/or wireless network connections that operate using a variety of network technologies. As a non-limiting example of a remote access controller, the integrated Dell Remote Access Controller (iDRAC) from Dell® is embedded within Dell servers and provides functionality that helps information technology (IT) administrators deploy, update, monitor, and maintain servers remotely.

(53) Remote access controller **230** supports monitoring and administration of the managed devices

of an IHS via a sideband bus 253. For instance, messages utilized in device and/or system management may be transmitted using I2C sideband bus 253 connections that may be individually established with each of the respective managed devices 205, 235a-b, 240, 250, 255, 260 of the IHS 200 through the operation of an I2C multiplexer 230d of the remote access controller. As illustrated in FIG. 2, the managed devices 205, 235a-b, 240, 250, 255, 260 of IHS 200 are coupled to the CPUs 205, either directly or indirectly, via in-line buses that are separate from the I2C sideband bus 253 connections used by the remote access controller 230 for device management.

(54) In certain embodiments, the service processor 230a of remote access controller 230 may rely on an I2C co-processor 230b to implement sideband I2C communications between the remote access controller 230 and the managed hardware components 205, 235a-b, 240, 250, 255, 260 of the IHS 200. The I2C co-processor 230b may be a specialized co-processor or micro-controller that is configured to implement an I2C bus interface used to support communications with managed hardware components 205, 235a-b, 240, 250, 255, 260 of HIS 200. In some embodiments, the I2C co-processor 230b may be an integrated circuit on the same die as the service processor 230a, such as a peripheral system-on-chip feature that may be provided by the service processor 230a. The I2C sideband bus 253 is illustrated as single line in FIG. 2. However, sideband bus 253 may be comprised of multiple signaling pathways, where each may be comprised of a clock line and data line that couple the remote access controller 230 to the I2C endpoints (e.g., managed hardware components 205, 235a-b, 240, 250, 255, 260).

(55) In various embodiments, an IHS 200 does not include each of the components shown in FIG. 2. In various embodiments, an IHS 200 may include various additional components in addition to those that are shown in FIG. 2. Furthermore, some components that are represented as separate components in FIG. 2 may in certain embodiments instead be integrated with other components. For example, in certain embodiments, all or a portion of the functionality provided by the illustrated components may instead be provided by components integrated into the one or more processor(s) 205 as a systems-on-a-chip.

(56) FIG. 3 is a diagram 300 illustrating several components of an IHS 200 showing how those components may communicate with one another for implementing an example rebootless firmware update system 300 according to one embodiment of the present disclosure. The IHS 200 is shown with a RAC software agent 302 and a basic input output system (BIOS) 304, and one or more device drivers 316 that are each associated with one of multiple firmware devices 306. The firmware devices 306 may be any IHS configurable device, such as a graphics processing card, a network interface card, an input/output (I/O) (e.g., USB) expansion port card, a Wi-Fi interface card, a storage controller card, and the like. For example, the firmware devices 306 may be a hardware accelerator 250, a storage drive 235a, an I/O port 215, a network controller 240, such as described herein above with reference to FIG. 2. Some, most, or all firmware devices 306 communicate with the IHS 200 via system bus 320, which in one embodiment, may include a Peripheral Component Interconnect Express (PCIe) bus.

(57) In one embodiment, the RAC 230 is configured to manage firmware updates for the various firmware devices 306 in the IHS 200. The RAC 230 may communicate with each of multiple firmware devices 306 via an i2c connection 308 such as described above to transfer the firmware update image 322 and activate it once transferred. In other embodiments, the RAC 230 may communicate with each of the firmware devices 306 using other communication protocols, such as i3c SENSEWIRE connections, or serial peripheral interface (SPI) based connections. Additionally, the RAC 230 may also communicate with each of the firmware devices 306 through the system bus 320 using a Management Component Transport Protocol (MCTP) PCIe vendor defined message (VDM) channel 310.

(58) The firmware devices 306 may be any type of IHS configurable device that may be updated with new firmware updates at an ongoing basis. For example, the firmware device 306 may include a non-volatile storage unit (e.g., hard disks, Solid State Drives (SSDs), etc.), Network Interface

Cards (NICs), Graphical Processing Units (GPUs), RACs, Hardware RAID (HWRAID) devices, and the like. Citing a particular example, the firmware device **306** may include a storage drive **235b**, those that are configured on a storage sled **115a-n**, and/or storage resources **155** configured in a JBOD, such as described herein above with reference to FIGS. **1** and **2**.

(59) The RAC **230** may also communicate with the IHS **200** through the PCIe VDM channel **310**, a Keyboard Controller Style (KCS) channel **312**, and/or a USB NIC channel **314**. The KCS channel **312** is provided as part of an intelligent platform management interface (IPMI) specification. In general, the IPMI specification defines several system interfaces for local access from the RAC software agent **302** and BIOS **304** to the RAC **230** of which the KCS protocol provides an ideal connection that does not require authentication for its operation.

(60) While the RAC **230** may communicate with each firmware device **306** for transferring its firmware update image **322**, and activating the firmware update image **322** following transfer, some communication links as described above may possess a higher throughput rate. Currently available firmware update images **322** may be up to 40 Mega-bytes in size, and when transferred over an I2C communication link, for example, the firmware update image **322** could consume a relatively large amount of time given that I2C communication links often require approximately 53 minutes to transfer a 100 Kilo-byte image. Other firmware devices using the same I2C communication link could potentially be starved for tens of minutes. Nevertheless, vendors of the firmware devices **306** often continue to use I2C communication links for download and activation of firmware update images **322** for various reasons, such as a reluctance to change what has previously worked in the past. Thus, a relatively large firmware update image **322** that requires a reboot of the IHS **200** could cause significant downtime to its overall operation.

(61) The RAC software agent **302** generally comprises a lightweight software service that is executed on the host IHS **200** to integrate certain operating system (OS) features with the RAC **230**. The RAC software agent **302** provides OS-related information to the RAC **230**, and may add capabilities such as LC log event replication into the OS log, WMI support (including storage), iDRAC SNMP alerts via OS, iDRAC hard reset and remote full Power Cycle. For example, the RAC software agent **302** may be an iDRAC Service Module (iSM) that is configured to operate with the integrated Dell remote access controller (iDRAC), which are both provided by DELL TECHNOLOGIES.

(62) The IHS **200** may receive a firmware update image **322** that is to be installed on multiple firmware devices **306**. The firmware update **322** typically includes multiple sub-images **324** that may each be associated with a particular component of the firmware device **306**. Each firmware update image **322** often possesses several sub-images **324** that are configured for updating certain portions of the overall firmware used by the firmware device **306**. A firmware update image **322** may, or may not include an updated version of each sub-image based upon whether it is needed or not. Reasons to update a version of a sub-image **324** may include, for example, to meet compliance policies, to take advantage of new features, enhancements to the IHS, deploy security fixes, and the like.

(63) Certain sub-images **324** may be critical to the updating of the firmware device **306** in that when updated, there exists a relatively high likelihood that the HIS **200** will need to be rebooted to activate the firmware update image **322** on the firmware device **306**. Within this disclosure, a critical sub-image **324** may include any component or portion of the firmware update image **322** that affects the core operation of the firmware update image **322** when installed on its respective firmware device **306**. Taken another way, the sub-images **324** may be any component or portion of the firmware update image **322** that may have a relatively high likelihood of requiring a reboot of the IHS **200** if it is updated on the firmware device **306**. Examples of such components may include portions that update the kernel in the firmware device **306**, system libraries that are used by other components in the firmware device **306**, bootstrapping firmware, and security stacks (e.g., OpenSSL, etc.) that maintain the integrity of the firmware device **306**.

(64) Embodiments of the present disclosure provide a solution to this problem, among others, by providing a rebootless firmware update system **300** that identifies the sub-images **324** in the firmware update image **322** that may be critical, and compares the versions of those sub-images **324** against the versions of the installed sub-components **326** of installed firmware **328** to determine if the critical sub-components **326** will be updated. If so, the rebootless firmware update system **300** may schedule a firmware update to be performed at a time in the future, such as at a time when the IHS **200** is incurring a minimal amount of processing load (e.g., during a maintenance window). In some embodiments, the rebootless firmware update system **300** may receive user input for providing manual selection of whether to continue to update the firmware device **306** immediately or schedule the firmware update for a later time.

(65) In one embodiment, the rebootless firmware update system **300** may be implemented via an Original Equipment Manufacturer (OEM) command (e.g., directive or instruction) that when issued by the RAC **230**, requests that the firmware device **306** respond with version information for each of its installed sub-components **326**. For example, when the RAC **230** receives a firmware update image **322** that is directed to a particular firmware device **306**, it may issue an OEM command **400** to that firmware device **306** to obtain information about the installed firmware sub-components **326** in the firmware device **306** so that a comparison may be made between the sub-images **324** and installed sub-components **326**.

(66) FIG. 4 illustrates an example of the various fields of an OEM command **400** that may be used to query the firmware device **306** for version information for each of its installed sub-components **326**. For example, the fields may include multiple request fields **402** indicating what information is requested by the RAC **230**, and multiple response fields **404** indicating how the requested information is to be included in a response OEM command **400**. Variable component data fields **406** indicate how the information about the sub-images **324** is to be parsed within the command.

(67) FIG. 5 illustrates an example sub-component version acquisition method **500** that may be used to acquire versioning information for the sub-components **326** of installed firmware **328** in the firmware devices **306** configured in an IHS **200** according to one embodiment of the present disclosure. In one embodiment, the sub-component version acquisition method **500** may be performed in whole, or in part, by the RAC **230** as described herein above. In other embodiments, the sub-component version acquisition method **500** may be performed by any suitable component or combination of components configured in the IHS **200**. In one embodiment, the sub-component version acquisition method **500** may be performed by a daemon that is executed as a background process to continually monitor for ongoing changes that may occur to the installed firmware **328** in each of the firmware devices **306**.

(68) At step **502**, the method **500** obtains an inventory of some, most, or all firmware devices **306** configured in the IHS **200**. For example, the method **500** may communicate with the BIOS **304** to obtain the inventory from its storage. At step **504**, the method **500** performs steps **506-510** for each firmware device **306** identified at step **502**. At step **506**, identifies the critical sub-components **326** in the installed firmware **328** of the firmware device **306**. For example, the method **500** may include logic that detects when a particular sub-component **326** is adapted to manipulate the kernel, system libraries, bootstrapping firmware, and/or any security stacks in the firmware device **306**. The method **500** may then record version information for each identified critical sub-component **326** at step **508**. For example, the method **500** may store information about each identified critical sub-component **326** in a memory of the RAC **230** and/or the IHS **200**. In one embodiment, the method **500** may use an OEM command such as described above with reference to FIG. 4 to obtain such information from the firmware device **306**.

(69) At step **510**, the method **500** determines whether all firmware devices **306** have been processed to obtain their sub-component **326** versioning information. If not, processing continues at step **504** to obtain versioning information for the sub-components **326** in another firmware device **306**; otherwise, processing continues at step **512** in which the method **500** waits for a rebootless

firmware update method **600** as described herein below. In other embodiments, the aforementioned method **500** may be performed any time that the rebootless firmware update system desires information about versioning information associated with the sub-components **326** of each firmware device **306** in the IHS **200**. For example, the method **500** may be performed each time the IHS **200** is rebooted, or each time a new firmware update image **322** is installed on one or more of the firmware devices **306** configured in the IHS **200**. For another example, the method **500** may be performed prior to rebootless firmware update method **600** being performed as will be described in detail herein below. Nevertheless, when use of the sub-component version acquisition method **500** is no longer needed or desired, the process ends.

(70) FIG. **6** illustrates an example rebootless firmware update method **600** depicting how a firmware device **306** configured in an IHS **200** may be updated according to one embodiment of the present disclosure. In one embodiment, the rebootless firmware update method **600** may be performed in whole, or in part, by the RAC **230** as described herein above. In other embodiments, the method **600** may be performed by any suitable component or combination of components that provide updates to the firmware device **306**. Initially, a new firmware update image **322** (e.g., software package or an updated version of an existing software package) is promoted or made available by a provider of the software package and/or the firmware device **306** that the firmware update image **322** supports.

(71) At step **602**, the method **600** receives a firmware update image **322**. For example, the method **600** may upload a firmware update image **322** into the RAC **230** so that the RAC **230** can administer the firmware update. At step **604**, the method **600** validates the firmware update image **322**. The RAC **230**, for example, may perform one or more validation steps to ensure that the downloaded firmware update image **322** is the correct type for the firmware device **306** to be updated, and/or any licensing agreements associated with the firmware device **306** are adhered to.

(72) The method **600** then determines whether the firmware update image **322** is validated at step **606**. If not, the method **600** at step **608** cancels the firmware update and notifies the user (e.g., IT administrator) that the firmware update image **322** has been invalidated in which the method **600** ends. If, however, the firmware update image **322** is validated, processing continues at step **610** in which the method obtains version information associated with the sub-images **324** of the firmware update image **322**. The firmware update image **322** may include a metadata portion that stores sub-image **324** version information for each of its constituent sub-images **324**.

(73) At step **612**, the method **600** obtains version information from the sub-components **326** of firmware **328** installed on firmware device **306**. In one embodiment, the version information may be obtained using the sub-component version acquisition method **500** as described above with reference to FIG. **5**, using instructions, such as a daemon that runs as a background process to continually update version information for each of the firmware devices **306** in the IHS **200**. In other embodiment, the method **600** may obtain the version information at any suitable time, such as while the rebootless firmware update method **600** is being performed.

(74) At step **614**, the method **600** compares the critical sub-images **324** in the firmware update image **322** with their corresponding installed sub-components **326** of the installed firmware **328**. Thereafter at step **616**, the method **600** determines whether any differences in the version exist between any critical sub-images **324** of the firmware update image **322** and their corresponding installed sub-components **326** of the installed firmware **328**. As mentioned previously, a difference in versions between the critical sub-images **324** in the firmware update image **322** with their corresponding installed sub-components **326** of the installed firmware **328** may indicate that the IHS **200** possesses a relatively high likelihood of requiring a reboot. If no any differences in the version exist between any critical sub-images **324** of the firmware update image **322** and their corresponding installed sub-components **326** of the installed firmware **328**, the firmware update of the firmware update image **322** may be scheduled at any suitable time (e.g., immediately) at step **618**. If, however, version differences of any of the critical sub-components **326** relative to the sub-

images 324 of the firmware update image 322 exist, processing continues at step 620 in which the firmware update is scheduled to be performed at a time in the future.

(75) In one embodiment, the method 600 may display an alert message (e.g., a popup window) for the user indicating that the firmware update of the firmware update image 322 will likely require a system reboot. In another embodiment, the method 600 may receive user input for determining whether to go ahead and schedule the firmware update of the firmware update image 322 immediately, or for selecting a particular time in the future to perform the firmware update. In another embodiment, the method 600 may include logic to schedule the firmware update based upon a known future time when a processing load is estimated to be at a minimal level, such as during a maintenance window. In yet another embodiment, the method 600 may restrict the firmware update from being performed. For example, the method 600 may restrict the firmware update from being performed by obtaining a processing load currently being incurred by the IHS 200, and determining that the current processing load exceeds a specified threshold.

(76) The aforescribed method 600 may be performed each time a firmware update image 322 is to be updated on one or more firmware devices 306 on the IHS 200 or other firmware devices 306 configured in a different IHS 200. Nevertheless, when use of the rebootless firmware update method 600 is no longer needed or desired, the process ends.

(77) Although FIGS. 5 and 6 describe example methods 500, 600 that may be performed to transfer firmware update images to firmware device(s) 306 in an IHS 200, the features of the disclosed processes may be embodied in other specific forms without deviating from the spirit and scope of the present disclosure. For example, certain steps of the disclosed methods 500, 600 may be performed sequentially, or alternatively, they may be performed concurrently. As another example, the methods 500, 600 may perform additional, fewer, or different operations than those operations as described in the present example.

(78) It should be understood that various operations described herein may be implemented in software executed by logic or processing circuitry, hardware, or a combination thereof. The order in which each operation of a given method is performed may be changed, and various operations may be added, reordered, combined, omitted, modified, etc. It is intended that the invention(s) described herein embrace all such modifications and changes and, accordingly, the above description should be regarded in an illustrative rather than a restrictive sense.

(79) Although the invention(s) is/are described herein with reference to specific embodiments, various modifications and changes can be made without departing from the scope of the present invention(s), as set forth in the claims below. Accordingly, the specification and figures are to be regarded in an illustrative rather than a restrictive sense, and all such modifications are intended to be included within the scope of the present invention(s). Any benefits, advantages, or solutions to problems that are described herein with regard to specific embodiments are not intended to be construed as a critical, required, or essential feature or element of any or all the claims.

(80) Unless stated otherwise, terms such as “first” and “second” are used to arbitrarily distinguish between the elements such terms describe. Thus, these terms are not necessarily intended to indicate temporal or other prioritization of such elements. The terms “coupled” or “operably coupled” are defined as connected, although not necessarily directly, and not necessarily mechanically. The terms “a” and “an” are defined as one or more unless stated otherwise. The terms “comprise” (and any form of comprise, such as “comprises” and “comprising”), “have” (and any form of have, such as “has” and “having”), “include” (and any form of include, such as “includes” and “including”) and “contain” (and any form of contain, such as “contains” and “containing”) are open-ended linking verbs. As a result, a system, device, or apparatus that “comprises,” “has,” “includes” or “contains” one or more elements possesses those one or more elements but is not limited to possessing only those one or more elements. Similarly, a method or process that “comprises,” “has,” “includes” or “contains” one or more operations possesses those one or more operations but is not limited to possessing only those one or more operations.

Claims

1. An Information Handling System (IHS) comprising: a firmware device that is configured to be updated with firmware on an ongoing basis; at least one processor; and a memory coupled to the at least one processor, the memory having program instructions stored thereon that, upon execution by the at least one processor, cause the IHS to: receive a firmware update image associated with the firmware device, the firmware update image conforming to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates, wherein the firmware update image comprises a plurality of sub-images and metadata for the plurality of sub-images; send a query to the firmware device, wherein the query is configured to cause the firmware device to respond with sub-component data associated with a plurality of firmware sub-components stored in the firmware device; receive the sub-component data from the firmware device, wherein the sub-component data comprises, for each firmware sub-component of the plurality of firmware sub-components stored in the firmware device: a sub-component functional identification that each sub-component is one of: a kernel, a system library, bootstrapping firmware, a security stack, or other component; a sub-component type indication that each firmware sub-component is one of: non-critical, critical, or unknown; and a sub-component version identification; determine, based at least in part on: the metadata included in the firmware update image, the sub-component functional identification received from the firmware device for each sub-component stored in the firmware device, the sub-component type indication, and the sub-component version identification received from the firmware device for each sub-component stored in the firmware device, that the firmware update image possesses a critical sub-image portion by detection that at least one of the plurality of sub-images is configured to perform at least one manipulation of a kernel, a system library, bootstrapping firmware, or a security stack in the firmware device, and that the at least one of the plurality of sub-images is a different version than the firmware sub-component that the at least one of the plurality of sub-images corresponds to; and in response to a determination that the firmware update image possesses the critical sub-image portion, increase IHS workload availability, comprising: schedule a firmware update with the firmware update image to be performed on the firmware device at a future time when a future processing load is estimated to be at a minimal level.
2. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to schedule the firmware update by receiving user input to specify the future time for the firmware update.
3. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to schedule the firmware update to be performed for all devices supported by the firmware update image in an IHS cluster comprising a plurality of IHSs.
4. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to generate an alert message for a user indicating that the firmware update of the firmware update image will likely require a system reboot.
5. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to: perform the firmware update at least in part via an integrated Dell Remote Access Controller (iDRAC).
6. The IHS of claim 5, wherein the instructions, upon execution, cause the IHS to perform the firmware update out-of-band.
7. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to receive the sub-component data using an Original Equipment Manufacturer (OEM) command that is configured to be issued to the firmware device.
8. The IHS of claim 1, wherein the instructions, upon execution, cause the IHS to receive the sub-component data via a daemon process.
9. The IHS of claim 1, wherein the instructions are performed by a Remote Access Controller (RAC) configured in the IHS.

10. A rebootless firmware update method comprising: receiving a firmware update image associated with a firmware device, the firmware update image conforming to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates, wherein the firmware update image comprises a plurality of sub-images and metadata for the plurality of sub-images; sending a query to the firmware device, wherein the query is configured to cause the firmware device to respond with sub-component data associated with a plurality of firmware sub-components stored in the firmware device; receiving the sub-component data from the firmware device, wherein the sub-component data comprises, for each firmware sub-component of the plurality of firmware sub-components stored in the firmware device: a sub-component functional identification that each sub-component is one of: a kernel, a system library, bootstrapping firmware, a security stack, or other component; and a sub-component type indication that each firmware sub-component is one of: non-critical, critical, or unknown; and a sub-component version identification; determining, based at least in part on: the metadata included in the firmware update image, the sub-component functional identification received from the firmware device for each sub-component stored in the firmware device, the sub-component type indication, and the sub-component version identification received from the firmware device for each sub-component stored in the firmware device, that the firmware update image possesses a critical sub-image portion by detecting that at least one of the plurality of sub-images is configured to perform at least one manipulation of a kernel, a system library, bootstrapping firmware, or a security stack in the firmware device, and that the at least one of the plurality of sub-images is a different version than the firmware sub-component that the at least one of the plurality of sub-images corresponds to; and in response to determining that the firmware update image possesses the critical sub-image portion, increasing IHS workload availability, comprising scheduling a firmware update with the firmware update image to be performed on the firmware device at a future time when a future processing load is estimated to be at a minimal level.

11. The rebootless firmware update method of claim 10, further comprising scheduling the firmware update by receiving user input to specify the future time for the firmware update.

12. The rebootless firmware update method of claim 10, further comprising generating an alert message for a user indicating that the firmware update of the firmware update image will likely require a system reboot.

13. The rebootless firmware update method of claim 10, further comprising receiving the sub-component data using an Original Equipment Manufacturer (OEM) command that is configured to be issued to the firmware device.

14. The rebootless firmware update method of claim 10, further comprising receiving the sub-component data using a daemon process.

15. A memory storage device having program instructions stored thereon that, upon execution by one or more processors of an Information Handling System (IHS), cause the IHS to: receive a firmware update image associated with a firmware device, the firmware update image conforming to a Platform Level Data Model (PLDM) protocol that is configured to perform rebootless firmware updates, wherein the firmware update image comprises a plurality of sub-images and metadata for the plurality of sub-images; send a query to the firmware device, wherein the query is configured to cause the firmware device to respond with sub-component data associated with a plurality of firmware sub-components stored in the firmware device; receive the sub-component data from the firmware device, wherein the sub-component data comprises, for each firmware sub-component of the plurality of firmware sub-components stored in the firmware device: a sub-component functional identification that each sub-component is one of: a kernel, a system library, bootstrapping firmware, a security stack, or other component; a sub-component type indication that each firmware sub-component is one of: non-critical, critical, or unknown; and a sub-component version identification; determine, based at least in part on: the metadata included in the firmware update image, the sub-component functional identification received from the firmware device for

each sub-component stored in the firmware device, the sub-component type indication, and the sub-component version identification received from the firmware device for each sub-component stored in the firmware device, that the firmware update image possesses a critical sub-image portion by detection that at least one of the plurality of sub-images is configured to perform at least one manipulation of a kernel, a system library, bootstrapping firmware, or a security stack in the firmware device, and that the at least one of the plurality of sub-images is a different version than the firmware sub-component that the at least one of the plurality of sub-images corresponds to; and in response to a determination that the firmware update image possesses the critical sub-image portion, increase IHS workload availability, comprising: schedule a firmware update with the firmware update image to be performed on the firmware device at a future time when a future processing load is estimated to be at a minimal level.

16. The memory storage device of claim 15, wherein the instructions, upon execution, cause the IHS to: schedule the firmware update to be performed for all devices supported by the firmware update image in an IHS cluster comprising a plurality of IHSs.

17. The memory storage device of claim 16, wherein the instructions, upon execution, cause the IHS to perform the firmware update at least in part via an integrated Dell Remote Access Controller (iDRAC).
