

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411694
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Wiencke; Christian et al.

Processor having adaptive pipeline with latency reduction logic that selectively executes instructions to reduce latency

Abstract

A system and method for reducing pipeline latency. In one embodiment, a processing system includes a processing pipeline. The processing pipeline includes a plurality of processing stages. Each stage is configured to further processing provided by a previous stage. A first of the stages is configured to perform a first function in a pipeline cycle. A second of the stages is disposed downstream of the first of the stages, and is configured to perform, in a pipeline cycle, a second function that is different from the first function. The first of the stages is further configured to selectably perform the first function and the second function in a pipeline cycle, and bypass the second of the stages.

Inventors: Wiencke; Christian (Garching, DE), Bhatia; Shrey Sudhir (Freising, DE), Vliegen; Jeroen (Fahrenzhausen, DE)

Applicant: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Family ID: 52481467

Assignee: TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)

Appl. No.: 18/314264

Filed: May 09, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20230273797 A1	Aug. 31, 2023

Related U.S. Application Data

continuation parent-doc US 13974571 20130823 US 11645083 child-doc US 18314264

Publication Classification

Int. Cl.: G06F9/38 (20180101)

U.S. Cl.:

CPC G06F9/3873 (20130101); G06F9/3838 (20130101); G06F9/3867 (20130101);

Field of Classification Search

CPC: G06F (9/3873); G06F (9/3875); G06F (9/3871); G06F (9/3867); G06F (9/3869); G06F (9/3826); G06F (9/383); G06F (9/3832)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5471626	12/1994	Carnevale	N/A	N/A
5619668	12/1996	Zaidi	N/A	N/A
5649174	12/1996	Dockser	N/A	N/A
5778250	12/1997	Dye	N/A	N/A
5802331	12/1997	van Berkel	N/A	N/A
6018796	12/1999	Suzuki	N/A	N/A
6839833	12/2004	Hartnett	N/A	N/A
8806181	12/2013	O'Bleness	N/A	N/A
9563432	12/2016	Segelken	N/A	N/A
11645083	12/2022	Wiencke	712/220	G06F 9/3838
2002/0199085	12/2001	Norden	N/A	N/A
2003/0208672	12/2002	Leenstra	N/A	N/A
2003/0226000	12/2002	Rhoades	N/A	N/A
2006/0095745	12/2005	Tran	N/A	N/A
2007/0271449	12/2006	Lichtensteiger	N/A	N/A
2008/0162891	12/2007	Pittman	N/A	N/A
2008/0244234	12/2007	Seth	N/A	N/A
2009/0138674	12/2008	Chang	N/A	N/A
2009/0198970	12/2008	Emma	N/A	N/A
2009/0292907	12/2008	Schwinn	N/A	N/A
2009/0327670	12/2008	Orbach	N/A	N/A
2014/0281423	12/2013	Caulfield	N/A	N/A

OTHER PUBLICATIONS

Stokes, J., et al., "Inside the Machine, An Illustrated Introduction to Microprocessors and Computer Architecture," No Starch Press, San Francisco, 7 pages. cited by applicant

Efthymiou, A., et al., "Adaptive Pipeline Structures for Speculation Control," Proceedings of the Ninth International Symposium on Asynchronous Circuits and Systems, The Computer Society, IEEE, 10 pages. cited by applicant

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) The present application is a continuation of U.S. patent application Ser. No. 13/974,571, filed Aug. 23, 2013, which is incorporated by reference herein in its entirety.

BACKGROUND

(1) Pipelining is one technique employed to increase the performance of processing systems such as microprocessors. Pipelining divides the execution of an instruction (or operation) into a number of stages where each stage corresponds to one step in the execution of the instruction. As each stage completes processing of a given instruction, and processing of the given instruction passes to a subsequent stage, the stage becomes available to commence processing of the next instruction. Thus, pipelining increases the overall rate at which instructions can be executed by partitioning execution into a plurality of steps that allow a new instruction to begin execution before execution of a previous instruction is complete. While pipelining increases the rate of instruction execution, pipelining also tends to increase instruction latency.

SUMMARY

(2) A system and method for reducing pipeline latency are disclosed herein. In one embodiment, a processor includes an execution pipeline and pipeline control logic. The execution pipeline includes a plurality of stages. The pipeline control logic is configured to identify an instruction being executed in the pipeline; to determine whether the identified instruction can be processed using fewer than a total number of the pipeline stages; and to selectably configure the pipeline to process the identified instruction using fewer than the total number of pipeline stages.

(3) In another embodiment, a processing system includes a processing pipeline. The processing pipeline includes a plurality of processing stages. Each stage is configured to further processing provided by a previous stage. A first of the stages is configured to perform a first function in a pipeline cycle. A second of the stages is disposed downstream of the first of the stages, and is configured to perform, in a pipeline cycle, a second function that is different from the first function. The first of the stages is further configured to selectably perform the first function and the second function in a pipeline cycle, and bypass the second of the stages.

(4) In a further embodiment, a method includes identifying, during execution, an instruction being executed in an execution pipeline comprising a plurality of stages. Whether the identified instruction can be processed using fewer than a total number of stages of the pipeline is determined. Responsive to the determination, the pipeline is configured to process the identified instruction using fewer than the total number of stages of the pipeline.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) For a detailed description of exemplary embodiments of the invention, reference will now be made to the accompanying drawings in which:

- (2) FIG. 1 shows a block diagram of a processing system in accordance with various embodiments;
- (3) FIG. 2 shows execution of instructions in a pipeline in accordance with various embodiments;
- (4) FIG. 3 shows a block diagram of a fetch unit in accordance with various embodiments;
- (5) FIG. 4 shows an alternative block diagram of a fetch unit in accordance with various embodiments; and

(6) FIG. 5 shows flow diagram for a method for executing instructions in an execution pipeline in accordance with various embodiments.

NOTATION AND NOMENCLATURE

(7) Certain terms are used throughout the following description and claims to refer to particular system components. As one skilled in the art will appreciate, companies may refer to a component by different names. This document does not intend to distinguish between components that differ in name but not function. In the following discussion and in the claims, the terms “including” and “comprising” are used in an open-ended fashion, and thus should be interpreted to mean “including, but not limited to . . .” Also, the term “couple” or “couples” is intended to mean either an indirect or direct electrical connection. Thus, if a first device couples to a second device, that connection may be through a direct electrical connection, or through an indirect electrical connection via other devices and connections. The recitation “based on” is intended to mean “based at least in part on.” Therefore, if X is based on Y, X may be based on Y and any number of additional factors.

DETAILED DESCRIPTION

(8) The following discussion is directed to various embodiments of the invention. Although one or more of these embodiments may be preferred, the embodiments disclosed should not be interpreted, or otherwise used, as limiting the scope of the disclosure, including the claims. In addition, one skilled in the art will understand that the following description has broad application, and the discussion of any embodiment is meant only to be exemplary of that embodiment, and not intended to intimate that the scope of the disclosure, including the claims, is limited to that embodiment.

(9) Processing systems, such as processors and other data manipulation systems, include pipelines to increase the processing rate of the system. The latency introduced by pipelining is generally considered acceptable given the increased throughput provided by the pipelining. However, in some situations pipeline latency can have a significant impact on system performance. Hazards, such as inter-instruction data dependencies (data hazards) and changes in instruction flow (control hazards) can result in an undesirable degree of additional pipeline latency. One conventional technique of dealing with pipeline hazards includes stalling the pipeline until the hazard is resolved. As pipeline length increases, the time the pipeline is stalled to resolve a hazard may also increase. Consequently, in conventional systems, pipeline stalls associated with hazard resolution can significantly degrade system performance.

(10) Embodiments of the processing system and pipeline disclosed herein reduce the latency associated with pipeline hazards, or other pipeline disruptions, and as a result, increase the overall throughput of the processing system. Embodiments reduce latency by varying the length of the pipeline based on the instruction or operation being executed in the pipeline. The pipeline identifies various operations that may cause a hazard, and reduces the length of the pipeline applied to execute the operation. By reducing the length of the pipeline applied to the operation, the number of pipeline cycles during which the pipeline is stalled in association with the operation is also reduced. For operations identified as not executable via a reduced length pipeline, embodiments may apply the full length pipeline.

(11) FIG. 1 shows a block diagram of a processing system **100** in accordance with various embodiments. The processing system **100** may be a processor, such as a general purpose microprocessor, a microcontroller, a digital signal processor, or other system that includes a processing pipeline. The system **100** includes a pipeline **102**. The pipeline **102** includes a plurality of successively coupled processing stages **104-110**. Various embodiments of the pipeline **102** may include more or fewer stages than are illustrated in FIG. 1.

(12) Each stage **104-110** provides processing functionality and each stage **106-110** provides processing functionality that furthers the processing provided by the previous stage. For example, in the pipeline **102**, stage 0 **104** may include a fetch unit that fetches instructions and/or data from

storage for execution/manipulation. Stage 1 **106** may include a decode unit that decodes instructions provided by the fetch unit of stage 0 **104**. Stage 2 **108** may include an execution unit that executes an instruction in accordance with the instruction decoding provided by stage 1 **106**. Stage 3 **110** may include a write-back unit that stores results of execution provided by the execution unit of stage 2 **108** to a selected storage device, such as memory or registers. The stages **104-110** may provide different functionality in some embodiments of the pipeline **102**.

(13) Stage 1 **106** also includes latency reduction logic **112**. The latency reduction logic **112** provides to stage 1 **106** functionality of succeeding stages **108** and/or **110** that is applied to execute functions of the succeeding stages in stage 1 **106** with respect to one or more selected instructions. For example, if stage 1 **106** is a decoding stage, then the latency reduction logic **112** may include execution logic used to execute selected instructions and writeback logic used to store the result of execution for the selected instructions. Thus, the pipeline **102** may execute the selected instructions in a reduced length pipeline that includes only stages **104** and **106**. In some embodiments of the pipeline **102**, latency reduction logic may be included in one or more stages of the pipeline **102** to provide various pipeline lengths in accordance with the instructions targeted for execution in each stage.

(14) In some embodiments of the system **100**, the latency reduction logic **112** may allow stage 1 **106** to perform the functions of succeeding pipeline stages with regard to instructions that can cause pipeline dependencies, and in turn cause pipeline hazards. In such embodiments, instructions that do not cause pipeline dependencies are executed using the full length of the pipeline rather than a reduced length pipeline. Such embodiments advantageously provide reduced pipeline latency when hazards occur, but maintain a high overall clock or execution rate by limiting the logic/functionality included in each pipeline stage.

(15) Pipeline control logic **114** is coupled to pipeline stage 1 **106** and the latency reduction logic **112**. The pipeline control logic **114** identifies the instructions being executed in the pipeline **102**, and selects, in accordance with the identified instruction, whether the pipeline stage **106** is to apply the latency reduction logic **112** to reduce pipeline length or to apply the full pipeline length. In some embodiments of the pipeline **102**, the pipeline control logic **114** may be included in stage **106**, e.g., in conjunction with or part of the latency reduction logic **112**.

(16) FIG. 2 shows execution of instructions in the pipeline **102** in accordance with various embodiments. More specifically, FIG. 2 shows execution of three instructions in the pipeline **102**. While executing Instruction 1, the pipeline control logic **114** identifies Instruction 1 as an instruction that cannot be executed in the reduced length pipeline formed of stages **104-106**. Accordingly, the pipeline control logic **114** configures stage 1 **106** to execute Instruction 1 without use of the latency reduction logic **112**, and Instruction 1 is executed using all the stages of the pipeline **102**.

(17) While executing Instruction 2, the pipeline control logic **114** identifies Instruction 2 as an instruction that can be executed in the reduced length pipeline formed of stages **104-106**. The pipeline control logic **114** may also evaluate an effect of execution of Instruction 2 on the pipeline **102**, and determine that the effect indicates that the Instruction 2 should be executed using the reduced length pipeline. Accordingly, the pipeline control logic **114** configures stage 1 **106** to execute Instruction 2 using the latency reduction logic **112**, and Instruction 2 is executed using the reduced length pipeline of stages **104-106**.

(18) The pipeline control logic **114** may identify Instruction 2 as causing a pipeline hazard. Execution of Instruction 2 in the reduced length pipeline reduces the latency caused by stalling the pipeline **102** to resolve the hazard. Execution of Instruction 2 using the reduced length pipeline stalls the execution of Instruction 3 by a single cycle, while execution of Instruction 2 using the full length pipeline **102** would have resulted in three stall cycles. Consequently, use of the reduced length pipeline to execute Instruction 2 allows Instruction 3 to be executed with less delay than had Instruction 2 been executed using the full length of the pipeline **102**. Thus, by reducing the length

of the pipeline applied to execute Instruction 2, pipeline latency is reduced, and performance of the system **100** is improved.

(19) In some embodiments of the pipeline **102**, the latency reduction logic **112** may perform only those operations of subsequent pipeline stages that are needed to reduce pipeline latency caused by execution of a selected instruction. Operations of the selected instruction execution not resulting in additional pipeline latency may be performed by the subsequent pipeline stages.

(20) In some embodiments of the pipeline **102**, a fetch stage of the pipeline **102** includes a fetch unit that includes an embodiment of the latency reduction logic **112** and the pipeline control logic **114**. In one such embodiment, the latency reduction logic **112** includes logic of each pipeline stage subsequent to the fetch stage for execution of program flow control instructions, such as jump, branch, call, etc., wholly in the fetch unit. FIG. **3** shows a block diagram of a fetch unit **300** including logic to execute flow control instructions in accordance with various embodiments. The fetch unit **300** includes fetch logic **302**, decode logic **304**, execution logic **306**, and writeback logic **308** applicable to execute the flow control instructions. The decode logic **304** can identify the flow control instructions. The execution logic **306** can determine the effect of an identified flow control instruction on the instruction stream. For example, the execution logic **306** may determine whether the identified flow control instruction redirects the instruction stream to a non-sequential instruction address, and determine the address of the next instruction to be executed. The writeback logic **308** can update a pointer to the next instruction to be executed. The fetch unit **300** may execute the operations of the logic **302-308** in a single pipeline cycle, or in fewer pipeline cycles than would be required to execute the equivalent operations using pipeline stages subsequent to the fetch stage.

(21) FIG. **4** shows an alternative block diagram of the fetch unit **300** in accordance with various embodiments. The block diagram of FIG. **4** shows the functionality provided by the logic **304-308**. The fetch unit **300** includes a program counter (PC) **402**, instruction identification logic **404**, instruction evaluation logic **406**, and PC update logic **408**. The PC **402** stores the address of the next instruction to be fetched and executed. The instruction identification logic **404** determines whether an instruction fetched is a flow control instruction. For example, the instruction identification logic **404** may compare opcodes of flow control instructions to the opcode of the current instruction.

(22) The instruction evaluation logic **406** determines whether execution of the instruction changes the address of the next instruction to be executed. For example, whether the current instruction is conditional may be determined, and if a condition code or other information needed to determine whether the program counter is to be nonsequentially updated is available, then the effect of execution of the instruction on the program address can be determined. If the instruction changes the address of the next instruction to be fetched, then the PC update logic **408** determines the address of the next instruction to be fetched, and provides the updated address to the PC **402**. The PC update logic **408** may include adders and other logic to modify the current PC based on an offset value, an address value, etc. provided with the instruction or otherwise stored in or available to the system **100** (e.g., stored in a general purpose register).

(23) FIG. **5** shows flow diagram for a method **500** for executing instructions in an execution pipeline in accordance with various embodiments. Though depicted sequentially as a matter of convenience, at least some of the actions shown can be performed in a different order and/or performed in parallel. Additionally, some embodiments may perform only some of the actions shown.

(24) In block **502**, the pipeline control logic **114** identifies the instruction being executed in pipeline stage **106**. Pipeline stage **106** includes latency reduction logic **112** that allows stage **106** to execute one or more selected instructions in the pipeline stage **106** without use of subsequent pipeline stages. The latency reduction logic **112** provides the functionality of subsequent stages needed to execute the one or more selected instructions without use of the subsequent stages.

- (25) The effect of execution of the instruction on the pipeline **102** may also be determined.
- (26) If, in block **504**, the instruction is identified as being an instruction that can be executed in a pipeline including a reduced number of stages, e.g., no stages subsequent to stage **106**, then in block **506** the pipeline stage **106** may be set to apply the functionality of the latency reduction logic **112** to execute the instruction. That is, stage **106** may be set to execute the instruction using a reduced length pipeline. Whether the instruction is to be executed in the reduced length pipeline may also be determined based on the determined effect of execution of the instruction. For example, if the instruction can be executed using fewer than all pipeline stages, but execution of the instruction using all pipeline stages does not detrimentally affect the pipeline (e.g., cause a hazard), then stage **106** may be set to execute the instruction without using the latency reduction logic **112**.
- (27) In block **508**, the instruction is executed using fewer than all the stages of the pipeline **102** (e.g., using no stages subsequent to stage **106**).
- (28) If, in block **504**, the instruction is identified as being an instruction that cannot be executed in a pipeline including a reduced number of stages, e.g., the latency reduction logic **112** lacks the functionality to execute the instruction without use of stages subsequent to stage **106**, then in block **510** the pipeline stage **106** may be set to apply single stage functionality to execute the instruction. That is, stage **106** may be set to execute the instruction using the full length pipeline, where stage **106** does not apply the functionality of the latency reduction logic **112**. Whether the instruction is to be executed in the full length pipeline may also be determined based on the determined effect of execution of the instruction. For example, if the instruction can be executed using fewer than all pipeline stages, but execution of the instruction using all pipeline stages does not cause a pipeline hazard, then the instruction may be executed using the full length pipeline.
- (29) In block **512**, the instruction is executed using all the stages of the pipeline **102**.
- (30) The above discussion is meant to be illustrative of the principles and various embodiments of the present invention. Numerous variations and modifications will become apparent to those skilled in the art once the above disclosure is fully appreciated. It is intended that the following claims be interpreted to embrace all such variations and modifications.

Claims

1. A method, comprising: identifying a first instruction being executed in a pipeline stage of an instruction processing pipeline; determining whether to perform an execute function on the first instruction in an execution stage or in a first stage, based on determining whether performing the execute function in the execution stage would cause a delay in execution of a second instruction that depends on the first instruction; and based on determining to perform the execute function on the first instruction in the first stage, stalling the instruction processing pipeline for a single pipeline cycle.
2. The method of claim 1, wherein the instruction processing pipeline comprises: a plurality of stages including: a fetch stage comprising circuitry configured to perform a fetch function on the first instruction; a decode stage coupled to the fetch stage and comprising circuitry configured to perform a decode function on the first instruction; and the execution stage, wherein the execution stage is coupled to the decode stage and comprises circuitry configured to perform the execute function on the first instruction.
3. The method of claim 2, further comprising: selecting the fetch stage or the decode stage to be the first stage.
4. The method of claim 2, wherein one of the fetch stage or the decode stage comprises a latency reduction circuit, the latency reduction circuit comprising circuitry that performs the execute function on the first instruction in a respective one of the fetch stage or the decode stage.
5. The method of claim 1, further comprising: based on determining to perform the execute function on the first instruction in the first stage, performing, by a latency reduction circuit separate

from the execution stage, the execute function on the first instruction in a decode stage.

6. The method of claim 1, further comprising: based on performing, by a latency reduction circuit, the execute function on the first instruction, storing, by the latency reduction circuit, a result of the execute function in a memory.

7. The method of claim 6, further comprising: determining whether the result is written to the memory by a writeback stage or by the latency reduction circuit.

8. The method of claim 1, further comprising: performing, by a latency reduction circuit, the execute function on the first instruction in a decode stage; and writing, by the latency reduction circuit, a result of the execute function produced in the decode stage.

9. A device comprising: a plurality of stages including: a fetch stage including circuitry configured to fetch a first instruction; a decode stage coupled to the fetch stage and including circuitry configured to decode the first instruction; and an execution stage coupled to the decode stage and including circuitry configured to execute the first instruction; a latency reduction circuit included within the decode stage, the latency reduction circuit including circuitry configured to execute the first instruction; and a pipeline control circuit coupled to the latency reduction circuit and configured to: determine whether to execute the first instruction by the execution stage or by the latency reduction circuit of the decode stage, based on determining whether executing the first instruction by the execution stage would cause a delay in execution of a second instruction that depends on the first instruction; and based on determining to execute the first instruction by the latency reduction circuit, stall operations of the plurality of stages for a single pipeline cycle.

10. The device of claim 9, wherein the latency reduction circuit further comprises circuitry configured to write a result of the execution of the first instruction to a memory.

11. The device of claim 9, wherein the plurality of stages includes a write-back stage, and wherein the pipeline control circuit is configured to: based on determining to execute the first instruction by the execution stage, cause the execution stage to execute the first instruction; and cause the write-back stage to write a result of the execution of the first instruction to a memory.

12. A method comprising: performing a first determination as to whether performing an execute operation on a first instruction using an execute stage delays a second instruction that depends on the first instruction; based on the first determination, performing a second determination as to whether to cause (a) the execute stage or (b) execute logic within a fetch stage or a decode stage to perform the execute operation on the first instruction; and based on the second determination, performing a third determination as to whether or not to stall the fetch stage, the decode stage, and the execute stage.

13. The method of claim 12, wherein: the execute logic performs a corresponding execute operation on each of a subset of a set of instructions; and the execute stage performs a corresponding execute operation on each of a remainder of the set of instructions.

14. The method of claim 13, wherein the subset of the set of instructions is a set of flow control instructions.

15. The method of claim 12, further comprising: determining whether to cause (a) a writeback stage or (b) writeback logic within the fetch stage or the decode stage to perform a writeback operation on the first instruction, wherein the writeback stage is coupled to the execute stage and is different from the writeback logic.

16. The method of claim 15, wherein the fetch stage includes: fetch logic that performs a fetch operation; decode logic that performs a decode operation; and the execute logic.

17. The method of claim 16, wherein the decode stage is coupled to the fetch stage and is different from the decode logic of the fetch stage, and performs the decode operation.

18. The method of claim 16, wherein the execute stage is coupled to the decode stage and is different from the execute logic.
