

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411712
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Vaka; Pratapa Reddy et al.

Predictable and adaptive quality of service for storage

Abstract

This disclosure describes a set of techniques that include establishing and managing quality of service standards across storage cluster resources in a data center. In one example, this disclosure describes a method that includes establishing a quality of service standard for a tenant sharing a storage resource with a plurality of tenants, wherein the storage resource is provided by the plurality of storage nodes in the storage cluster; allocating a volume of storage within the storage cluster, wherein allocating the volume of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes; and scheduling operations to be performed by the set of storage nodes for the volume of storage.

Inventors: Vaka; Pratapa Reddy (Saratoga, CA), Sindhu; Pradeep (Los Altos Hills, CA), Kohli; Jaspal (Sunnyvale, CA), Vegesna; Srihari Raju (San Jose, CA)

Applicant: Microsoft Technology Licensing, LLC (Redmond, WA)

Family ID: 86500154

Assignee: Microsoft Technology Licensing, LLC (Redmond, WA)

Appl. No.: 17/457230

Filed: December 01, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20230168934 A1	Jun. 01, 2023

Publication Classification

Int. Cl.: G06F9/50 (20060101); G06F9/48 (20060101); H04L67/1097 (20220101)

U.S. Cl.:

CPC **G06F9/5022** (20130101); **G06F9/4881** (20130101); **H04L67/1097** (20130101);
G06F2209/505 (20130101)

Field of Classification Search

CPC: G06F (9/5022); G06F (9/4881); G06F (2209/505); G06F (2209/5016); H04L (67/1097)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
9176708	12/2014	Gallant et al.	N/A	N/A
9515956	12/2015	Hu et al.	N/A	N/A
10481801	12/2018	Baitule et al.	N/A	N/A
10540288	12/2019	Noureddine et al.	N/A	N/A
10565112	12/2019	Noureddine et al.	N/A	N/A
10574579	12/2019	Puttagunta et al.	N/A	N/A
10659254	12/2019	Sindhu et al.	N/A	N/A
10686729	12/2019	Sindhu et al.	N/A	N/A
10841245	12/2019	Gray et al.	N/A	N/A
10929175	12/2020	Goyal et al.	N/A	N/A
11048634	12/2020	Noureddine et al.	N/A	N/A
11178262	12/2020	Goel et al.	N/A	N/A
2003/0079018	12/2002	Lolayekar	709/239	G06F 3/067
2003/0079019	12/2002	Lolayekar	709/226	H04L 67/101
2003/0191857	12/2002	Terrell	709/244	H04L 49/90
2003/0223445	12/2002	Lodha	370/428	H04L 47/10
2004/0133622	12/2003	Clubb	709/200	G06F 9/5077
2004/0193827	12/2003	Mogi	711/170	G06F 11/3433
2005/0235072	12/2004	Smith	710/22	G06F 13/28
2006/0106749	12/2005	Ganfield	N/A	G06F 16/9024
2007/0015525	12/2006	Beming	370/235	H04L 47/83
2010/0131957	12/2009	Kami	718/104	G06F 9/5077
2012/0317358	12/2011	Ando	711/E12.002	G06F 3/067
2014/0337577	12/2013	Burton	711/114	G06F 3/0631
2019/0108231	12/2018	Patterson	N/A	G06F 11/1464
2019/0236150	12/2018	Zaslavsky	N/A	G06F 16/2308
2020/0167097	12/2019	Huen et al.	N/A	N/A
2020/0183841	12/2019	Noureddine et al.	N/A	N/A
2020/0280462	12/2019	Sindhu et al.	N/A	N/A
2020/0314026	12/2019	Sindhu et al.	N/A	N/A
2021/0349824	12/2020	Noureddine et al.	N/A	N/A

OTHER PUBLICATIONS

U.S. Appl. No. 17/454,731, filed Nov. 12, 2021, naming inventors Goel et al. cited by applicant

Background/Summary

TECHNICAL FIELD

(1) This disclosure relates to sharing resources in the fields of networking and data storage.

BACKGROUND

(2) With advances in data center fabric technology, storage capacity, and networking speeds, storage systems in data centers are evolving. A storage cluster is a system enabling efficient storage of data within a data center or across data centers, and enabling access to such data to customers or tenants of a data center that share the resources of the storage cluster. Because there might be many tenants sharing resources of a data center, customer service level agreements (SLAs) are sometimes used to establish quality of service (QoS) standards. Such standards may help ensure not only that each tenant receives an expected or agreed-upon level of service, but such standards may also reduce adverse effects of noisy neighbor tenants in a data center that might disrupt other tenants' use of the data center.

SUMMARY

(3) This disclosure describes techniques that include establishing and managing quality of service (QoS) standards across storage cluster resources in a data center. In some examples, such techniques may involve establishing quality of service standards for customers, tenants, and/or operations across multiple storage cluster resources and/or multiple computing systems. To effectively manage such QoS standards, an orchestration system within the storage cluster may allocate, in a hierarchical way, storage cluster resources. Further, computing devices or computing nodes within the storage cluster may collectively schedule operations to be performed using the resources within the storage cluster. Scheduling operations may involve applying an algorithm that seeks to ensure a guaranteed availability of resources associated with a given storage unit (e.g., a "volume" of storage) within the storage cluster and also a maximum availability of resources associated with the storage unit. Such guaranteed and maximum levels of service may apply to multiple types of resources (e.g., storage capacity, processing cycles, bandwidth, and others) as well as to multiple operations associated with a resource (e.g., read and write operations).

(4) Techniques described herein may provide certain technical advantages. For instance, by taking QoS standards into account when allocating storage cluster resources, dynamically moving storage units (e.g., volumes) when needed, and limiting (e.g., rate limiting) use of resources within the storage cluster, it is possible to efficiently use a diverse set of resources that perform multiple types of operations across the storage cluster.

(5) In some examples, this disclosure describes operations performed by a compute node, storage node, computing system, network device, and/or storage cluster in accordance with one or more aspects of this disclosure. In one specific example, this disclosure describes a method comprising establishing, by a storage cluster having a plurality of storage nodes, a quality of service standard for a tenant sharing a storage resource with a plurality of tenants, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, and wherein the storage resource is provided by the plurality of storage nodes in the storage cluster; allocating, by the storage cluster and based on the quality of service standard, a volume of storage within the storage cluster, wherein allocating the volume of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage

nodes; and scheduling, by the storage cluster and based on the quality of service standard, operations to be performed by the set of storage nodes for the volume of storage.

(6) In another example, this disclosure describes a storage cluster comprising: a network; and a plurality of computing systems, each interconnected over the network and including a plurality of storage nodes, wherein the plurality of computing systems are collectively configured to: establish a quality of service standard for a tenant sharing a storage resource with a plurality of tenants, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, and wherein the storage resource is provided by the plurality of storage nodes in the storage cluster, allocate, based on the quality of service standard, a volume of storage within the storage cluster, wherein allocating the volume of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes, and schedule, based on the quality of service standard, operations to be performed by the set of storage nodes for the volume of storage.

(7) In another example, this disclosure describes a storage cluster comprising processing circuitry and a system for storing computing instructions, wherein the processing circuitry has access to the system for storing computing instructions and is configured to: establish a quality of service standard for a tenant sharing a storage resource with a plurality of tenants, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, and wherein the storage resource is provided by the plurality of storage nodes in the storage cluster, allocate, based on the quality of service standard, a volume of storage within the storage cluster, wherein allocating the volume of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes, and schedule, based on the quality of service standard, operations to be performed by the set of storage nodes for the volume of storage.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1A is a block diagram illustrating an example system including one or more network devices configured to efficiently process and store data reliably in a storage cluster, in accordance with one or more aspects of the present disclosure.
- (2) FIG. 1B is a simplified block diagram illustrating an example storage cluster, in accordance with one or more aspects of the present disclosure.
- (3) FIG. 2A is a block diagram illustrating an example node or data processing unit (DPU), in accordance with the techniques of this disclosure.
- (4) FIG. 2B is a conceptual diagram illustrating one possible partitioning of virtual processors within an example processor, in accordance with one or more aspects of the present disclosure.
- (5) FIG. 3 is a block diagram illustrating another example of an example DPU, in accordance with one or more aspects of the present disclosure.
- (6) FIG. 4 is a conceptual diagram illustrating various roles played by one or more nodes in a storage cluster, in accordance with one or more aspects of the present disclosure.
- (7) FIG. 5A is a conceptual diagram illustrating a model for scheduling a single resource within an example storage cluster, in accordance with one or more aspects of the present disclosure.
- (8) FIG. 5B illustrates a storage data path distributed across physically distinct nodes or data processing units, in accordance with one or more aspects of the present disclosure.
- (9) FIG. 6 is a flow diagram illustrating operations performed by an example storage cluster in accordance with one or more aspects of the present disclosure.

DETAILED DESCRIPTION

(10) FIG. 1A is a block diagram illustrating an example system **108** including one or more network devices configured to efficiently process and store data reliably in a storage cluster, in accordance with one or more aspects of the present disclosure. System **108** may include or represent a data center capable of performing data storage operations pursuant to quality of service (QoS) standards and/or service level agreements. Techniques described herein may enable efficient and effective compliance with such standards and/or agreements. Nodes as described herein may also be referred to as data processing units (DPUs) or devices including DPUs. For example, in FIG. 1A, various processing techniques are performed by nodes **17** within data center **101**. Other devices within a network, such as routers, switches, servers, firewalls, gateways and the like, may readily be configured to utilize the data processing techniques described herein.

(11) Data center **101** represents an example of a system in which various techniques described herein may be implemented. In general, data center **101** provides an operating environment for applications and services for tenants or customers **11** coupled to the data center **101** by service provider network **7** and gateway device **20**. Data center **101** may, for example, host infrastructure equipment, such as compute nodes, networking and storage systems, redundant power supplies, and environmental controls. Service provider network **7** may be coupled to one or more networks administered by other providers, and may thus form part of a large-scale public network infrastructure, e.g., the Internet.

(12) In some examples, data center **101** may represent one of many geographically distributed network data centers. In the example of FIG. 1A, data center **101** is a facility that provides information services for tenants **11**. Tenants or customers **11** may be collective entities such as enterprises and governments or individuals. For example, a network data center may host web services for several enterprises and end users. Other exemplary services may include data storage, virtual private networks, file storage services, data mining services, scientific- or super-computing services, and so on.

(13) Controller **130**, shown included within data center **101** of FIG. 1A, may be one or more computing devices that manage aspects of how data center **101** is configured and/or operates. In some examples, controller **130** may operate as a high-level controller or may serve as a software-defined networking (SDN) controller that configures and manages the routing and switching infrastructure of data center **101**. In such an example, controller **130** may provide a logically (and in some cases physically) centralized controller for facilitating operation of one or more virtual networks within data center **101**. Controller **130** may operate on its own, or in response to signals received from an administrator device (not shown) operated by an administrator. Controller **130** may offer application programming interface (“API”) support for various cluster services, which may include orchestration, storage services, and/or storage management capabilities. Such capabilities may also include infrastructure discovery, registration, and initialization, role-based access control, multi-tenancy and resource partitioning, application workload deployment and orchestration, flexible network control, identity management, and hardware lifecycle management and monitoring.

(14) Controller **130** may also be responsible for allocating and accounting for resources for a “volume,” which may, in some examples, refer to a basic storage unit abstraction supported by a data center or a storage cluster within a data center. In such an example, a volume may be a storage container divided into fixed size blocks, and be capable of being allocated and deallocated by controller **130**, as well as being written to and read from by nodes or other devices within the data center.

(15) In the illustrated example, data center **101** includes a set of storage systems and application servers **12** interconnected via a high-speed switch fabric **114**. In some examples, servers **12** are arranged into multiple different server groups, each including any number of servers up to, for example, n servers **12.sub.1-12N**. Servers **12** provide computation and storage facilities for

applications and data associated with tenants or customers **11** and may be physical (bare-metal) servers, virtual machines running on physical servers, virtualized containers running on physical servers, or combinations thereof.

(16) In the example of FIG. **1A**, some of servers **12** may be coupled to switch fabric **114** by one or more nodes **17** for processing streams of information, such as network packets or storage packets. In example implementations, nodes **17** may be configurable to operate in a standalone network appliance having one or more nodes. For example, nodes **17** may be arranged into multiple different node groups **19**, each including any number of nodes up to, for example, “N” nodes **17.sub.1-17.sub.N** (representing any number of nodes **17**). In other examples, each node may be implemented as a component (e.g., electronic chip) within a device, such as a compute node, application server, storage server, and may be deployed on a motherboard of the device or within a removable card, such as a storage and/or network interface card.

(17) In the example shown in FIG. **1A**, some nodes **17** are shown connected to one or more servers **12**, and such nodes **17** may serve to offload (from servers **12**) aspects of the handling of data packets and other network-related functions. These nodes are shown logically or physically organized within node groups **19**, units **40**, and racks **70**. Specifically, rack **70-1** includes one or more node groups **19**, each including a set of nodes **17** and storage devices **127**. The node group and the set of servers **12** supported by the nodes **17** of the node group **19** may be referred to herein as a network storage compute unit (NCSU) **40**. Illustrated in FIG. **1A** are NCSU **40-1** through NCSU **40-N**, which represent any number of NCSUs. (For ease of illustration, only components of NCSU **40-1** are shown.) In some examples, data center **101** may include many NCSUs, and multiple NCSUs **40** may be organized into logical racks or physical racks within data center **101**. For example, in some implementations, two NCSUs may compose a logical rack, and four NCSUs may compose a physical rack **70-1**. Other arrangements are possible. Such other arrangements may include nodes **17** within rack **70-1** being relatively independent, and not logically or physically included within any node group or NCSUs **40**.

(18) In general, each node group **19** of rack **70-1** may be configured to operate as a high-performance I/O hub designed to aggregate and process network and/or storage I/O for multiple servers **12**. As mentioned above, the set of nodes **17** within each of the node groups **19** provide programmable, specialized I/O processing circuits for handling networking and communications operations on behalf of servers **12**. In addition, in some examples, each of node groups **19** may include storage devices **127**, such as solid state drives (SSDs) and/or hard disk drives (HDDs), configured to provide network accessible storage for use by applications executing on the servers **12**. In some examples, one or more of the SSDs may comprise non-volatile memory (NVM) or flash memory. Although illustrated as logically within node groups **19** and external to nodes **17**, storage devices may alternatively or in addition be included within one or more nodes **17** or within one or more servers **12**.

(19) Other nodes **17** may serve as storage nodes (“storage targets”) that might not be directly connected to any of servers **12**. For instance, FIG. **1A** illustrates rack **70-2**, which includes nodes **17-1** through **17-N** (representing any number of nodes **17**). These nodes **17** may be configured to store data within one or more storage devices **127** (included within or connected to such nodes **17**) in accordance with techniques described herein. In the example illustrated, nodes **17** within rack **70-2** are not organized into groups or units, but instead, are relatively independent of each other, and are each capable of performing storage functions described herein. In other examples, however, nodes **17** of rack **70-2** may be logically or physically organized into groups, units, and/or logical racks as appropriate.

(20) Rack **70-3** is illustrated as being implemented in a manner similar to rack **70-2**, with storage nodes **17** configured to store data within storage devices. Although for ease of illustration, only racks **70-1**, **70-2**, **70-3**, through **70-N** are illustrated or represented in FIG. **1A**, any number of racks **70** may be included within data center **101**. Further, although rack **70-1** in FIG. **1A** is illustrated

with nodes **17** that support servers **12** and other racks **70** are illustrated with nodes **17** serving as storage nodes, in other examples, any number of racks may include nodes **17** that support servers **12**, and any number of racks may include nodes serving as storage nodes. Further, any of racks **70** may include a mix of nodes **17** supporting servers **12** and nodes **17** serving as storage nodes. Still further, although data center **101** is illustrated in the context of nodes **17** being arranged within racks **70**, other logical or physical arrangements of nodes **17** may be used in other implementations, and such other implementation may involve groups, units, or other logical or physical arrangements not involving racks.

(21) Nodes **17** of rack **70-2** (or rack **70-3**) may be devices or systems that are the same as or similar to nodes **17** of rack **70-1**. In other examples, nodes **17** of rack **70-2** may have different capabilities than those of rack **70-1** and/or may be implemented differently. In particular, nodes **17** of rack **70-2** may be somewhat more capable than nodes **17** of rack **70-1**, and may have more computing power, more memory capacity, more storage capacity, and/or additional capabilities. For instance, each of nodes **17** of rack **70-2** may be implemented by using a pair of nodes **17** of rack **70-1**. To reflect such an example, nodes **17** of rack **70-2** and **70-3** are illustrated in FIG. **1A** as being larger than nodes **17** of rack **70-1**.

(22) In a large scale fabric, storage systems (e.g., represented by nodes **17** of rack **70-2** or even NCSUs **40** of rack **70-1**) may become unavailable from time to time. Failure rates of storage systems are often significant, even if single component failure rates are quite small. Further, storage systems may become unavailable for reasons other than a software error or hardware malfunction, such as when a storage system or other device is being maintained or the software on such a device is being modified or upgraded. Data durability procedures may be employed to ensure access to critical data stored on a network when one or more storage systems are unavailable.

(23) In some examples, one or more hardware or software subsystems may serve as a failure domain or fault domain for storing data across data center **101**. For instance, in some examples, a failure domain may be chosen to include hardware or software subsystems within data center **101** that are relatively independent, such that a failure (or unavailability) of one such subsystem is relatively unlikely to be correlated with a failure of another such subsystem. Storing data fragments in different failure domains may therefore reduce the likelihood that more than one data fragment will be lost or unavailable at the same time. In some examples, a failure domain may be chosen at the node level, where each node represents a different failure domain. In another example, a failure domain may be chosen at a logical or physical grouping level, such that each group or unit of nodes **17** represents a different failure domain. In other examples, failure domains may be chosen more broadly, so that a failure domain encompasses a logical or physical rack **70** comprising many nodes **17**. Broader or narrower definitions of a failure domain may also be appropriate in various examples, depending on the nature of the network **8**, data center **101**, or subsystems within data center **101**.

(24) As further described herein, in one example, each node **17** may be a highly programmable I/O processor specially designed for performing storage functions and/or for offloading certain functions from servers **12**. In one example, each node **17** includes a number of internal processor clusters, each including two or more processing cores and equipped with hardware engines that offload cryptographic functions, compression and regular expression processing, data durability functions, data storage functions and networking operations. In such an example, each node **17** may include components for processing and storing network data (e.g., nodes **17** of rack **70-2**) and/or for and processing network data on behalf of one or more servers **12** (e.g., nodes **17** of rack **70-1**). In addition, some or all of nodes **17** may be programmatically configured to serve as a security gateway for its respective servers **12**, freeing up other computing devices (e.g., the processors of the servers **12**) to dedicate resources to application workloads.

(25) In various example implementations, some nodes **17** may be viewed as network interface subsystems that serve as a data storage node configured to store data across storage devices **127**.

Other nodes **17** in such implementations may be viewed as performing full offload of the handling of data packets (with, in some examples, zero copy in server memory) and various data processing acceleration for the attached server systems.

(26) In one example, each node **17** may be implemented as one or more application-specific integrated circuit (ASIC) or other hardware and software components, each supporting a subset of storage devices **127** or a subset of servers **12**. In accordance with the techniques of this disclosure, any or all of nodes **17** may include a data durability module or unit, which may be implemented as a dedicated module or unit for efficiently and/or quickly performing data durability operations. In some examples, such a module or unit may be referred to as an “accelerator” unit. That is, one or more computing devices may include a node including one or more data durability, data reliability, and/or erasure coding accelerator units.

(27) In the example of FIG. **1A**, each node **17** provides storage services (e.g., nodes **17** of rack **70-2**) or connectivity to switch fabric **114** for a different group of servers **12** (e.g., nodes **17** of rack **70-1**). Each of nodes **17** may be assigned respective IP addresses and provide routing operations for servers or storage devices coupled thereto. Nodes **17** may interface with and utilize switch fabric **114** so as to provide full mesh (any-to-any) interconnectivity such that any nodes **17** (or servers **12**) may communicate packet data for a given packet flow to any node **17** using any of a number of parallel data paths within the data center **101**. In addition, nodes **17** described herein may provide additional services, such as security (e.g., encryption), acceleration (e.g., compression), data reliability (e.g., erasure coding), I/O offloading, and the like. In some examples, each of nodes **17** may include or have access to storage devices, such as high-speed solid-state drives or rotating hard drives, configured to provide network accessible storage for use by applications executing on the servers. More details on the data center network architecture and interconnected nodes illustrated in FIG. **1A** are available in U.S. Pat. No. 10,686,729, entitled “Non-Blocking Any-to-Any Data Center Network with Packet Spraying Over Multiple Alternate Data Paths,” the entire content of which is incorporated herein by reference.

(28) Example architectures of nodes **17** are described herein with respect to FIG. **2A** and FIG. **3**. For some or all of such examples, the architecture of each node **17** comprises a multiple core processor system that represents a high performance, hyper-converged network, storage, and data processor and input/output hub. The architecture of each node **17** may be optimized for high performance and high efficiency stream processing. For purposes of example, DPUs corresponding to or within each node **17** may execute an operating system, such as a general-purpose operating system (e.g., Linux or Unix) or a special-purpose operating system, that provides an execution environment for data plane software for data processing.

(29) More details on how nodes **17** may operate are available in U.S. Pat. No. 10,841,245, entitled “Work Unit Stack Data Structures in Multiple Core Processor System,” U.S. Pat. No. 10,540,288, entitled “EFFICIENT WORK UNIT PROCESSING IN A MULTICORE SYSTEM”, filed Feb. 2, 2018, and in U.S. Pat. No. 10,659,254, entitled “Access Node Integrated Circuit for Data Centers which Includes a Networking Unit, a Plurality of Host Units, Processing Clusters, a Data Network Fabric, and a Control Network Fabric.” All of these publications are hereby incorporated by reference.

(30) FIG. **1B** is a simplified block diagram illustrating an example storage cluster, in accordance with one or more aspects of the present disclosure. FIG. **1B** illustrates storage cluster **102**, which may be considered to be an example storage cluster included within data center **101** of FIG. **1A**. Storage cluster **102** in FIG. **1B** is similar to the illustration of data center **101** of FIG. **1A**, and includes many of the same components illustrated in FIG. **1A**. However, elements of FIG. **1A** have been rearranged within FIG. **1B** to help illustrate certain aspects of how storage cluster **102** might be implemented within data center **101**.

(31) In the example of FIG. **1B**, storage cluster **102** includes controller **130**, one or more initiator nodes **110**, and one or more storage nodes **120**, all capable of communicating through switch fabric

114. One or more volumes **121** (e.g., volume **121J** and volume **121K**) each represent a “volume,” which might be considered a conceptual abstraction of a unit of storage in storage cluster **102**. Volumes may be associated with different tenants or customers of data center **101** of storage cluster **102**. For example, in the example illustrated in FIG. **1B**, volume **121J** has been allocated for use by tenant J, while volume **121K** has been allocated for use by tenant K. In FIG. **1B**, dotted lines radiating from each of volumes **121J** and **121K** are intended to illustrate that such volumes **121** are each stored across multiple storage nodes **120**. Although only two volumes are illustrated in FIG. **1B**, storage cluster **102** may support many more volumes **121** for many more tenants.

(32) As in FIG. **1A**, controller **130** provides cluster management orchestration of storage resources within storage cluster **102**. Also, as in FIG. **1A**, controller **130** may be implemented through any suitable computing system, including one or more compute nodes within data center **101** or storage cluster **102**. Although illustrated as a single system within storage cluster **102** in FIG. **1B**, controller **130** may be implemented as multiple system and/or as a distributed system that resides both inside and outside data center **101** and/or storage cluster **102**. In other examples, controller **130** some or all aspects of may be implemented outside of data center **101**, such as in a cloud-based implementation.

(33) In the example shown, controller **130** includes storage services module **131** and data store **132**. Storage services module **131** of controller **130** may perform functions relating to establishing, allocating, and enabling read and write access to one or more volumes **121** within storage cluster **102**. In general, storage services module **131** may perform functions that can be characterized as “cluster services,” which may include allocating, deleting, creating, and/or deleting volumes. As described herein, storage services module **131** may also provide services that help ensure compliance with quality of service standards for volumes **121** within storage cluster **102**. In some examples, storage services module **131** may also manage input from one or more administrators (e.g., operating administrator device **133**). In general, storage services module **131** may have a full view of all resources within storage cluster **102** and how such resources are allocated across volumes **121**.

(34) Data store **132** may represent any suitable data structure or storage medium for storing information related to resources within storage cluster **102**, and how such resources are allocated within storage cluster **102** and/or across volumes **121**. Data store **132** may be primarily maintained by storage services module **131**.

(35) Each of initiator nodes **110** may correspond to or be implemented by one or more of the servers **12** and nodes **17** illustrated in FIG. **1A**. Specifically, each of initiator nodes **110** is shown in FIG. **1B** as including at least one server **112** and DPU **118**. Each server **112** within initiator nodes **110** of FIG. **1B** may correspond to one or more of servers **12** of FIG. **1A**. Similarly, each DPU **118** within initiator nodes **110** of FIG. **1B** may correspond to one or more of nodes **17** (or DPUs **17**) of FIG. **1A**. The descriptions of servers **12** and nodes **17** provided in connection with FIG. **1A** may therefore apply to servers **112** and DPUs **118** of FIG. **1B**.

(36) Initiator nodes **110** illustrated in FIG. **1B** may be involved in causing or initiating a read and/or write operation with the storage cluster represented by storage cluster **102**. DPUs **118** within each of initiator nodes **110** may serve as the data-path hub for each of initiator nodes **110**, connecting each of initiator nodes **110** (and storage nodes **120**) through switch fabric **114**. In some examples, one or more of initiator nodes **110** may be an x86 server that may execute NVMe (Non-Volatile Memory Express) over a communication protocol, such as TCP. In some examples, other protocols may be used, including, for example, “FCP” as described in United States Patent Publication No. 2019-0104206 A1, entitled “FABRIC CONTROL PROTOCOL FOR DATA CENTER NETWORKS WITH PACKET SPRAYING OVER MULTIPLE ALTERNATE DATA PATH,” and which is hereby incorporated by reference.

(37) Each of storage nodes **120** may be implemented by the nodes **17** and storage devices **127** that are illustrated in FIG. **1A**. Accordingly, the description of such nodes **17** and storage devices **127** in

FIG. 1A may therefore apply to DPUs 117 and storage devices 127 of FIG. 1B, respectively. Storage nodes 120 are illustrated in FIG. 1B to emphasize that in some examples, each of storage nodes 120 may serve as storage targets for initiator nodes 110 in FIG. 1B.

(38) FIG. 1B also includes conceptual illustrations of volumes 121J and 121K. Within storage cluster 102, volumes 121 may serve as storage containers for data associated with tenants of storage cluster 102, where each such volume is an abstraction intended to represent a set of data that is stored across one or more storage nodes 120 of FIG. 1B. In some examples, each of volumes 121 may be divided into fixed size blocks and may support multiple operations. Typically, such operations generally include a read operation (i.e., reading one or more fixed-size blocks from a volume) and a write operation (i.e., writing one or more fixed-size blocks to a volume). Other operations are possible and are within the scope of this disclosure.

(39) Often, numerous tenants share resources of storage cluster 102, including storage resources. To communicate or indicate the level of service a current or prospective tenant can expect from storage cluster 102, a service level agreement (“SLA”) may be established between the operator of storage cluster 102 and a tenant or customer seeking to use services provided by storage cluster 102. Such SLAs may specify quality of service (QoS) standards that are used not only to ensure that each tenant gets the expected level of service (e.g., a “guaranteed service level”), but also to avoid a “noisy neighbor” problem arising from one tenant using so many resources of storage cluster 102 that such use disrupts or impacts the services provided to other tenants. Metrics that can be evaluated in order to assess or establish a QoS in a storage cluster might include processing operations and/or bandwidth measured in input/output operations per second (“IOPs”) and latency measured in microseconds.

(40) As described herein, a quality of service standard may include a guaranteed level of service. This may mean that resources needed for a storage service offered to a tenant should always be available from storage cluster 102 when needed. Storage cluster 102 may ensure that such guaranteed levels of service are met by managing and provisioning resources within storage cluster 102 (e.g., DPUs 117, storage devices 127, network resources, bandwidth, as well as others). Storage cluster 102 may also ensure that such guaranteed levels of service are met by appropriately allocating, placing, moving volumes within storage cluster 102, and in addition, rate limiting various operations involving the volumes.

(41) In addition, a quality of service standard may enable tenants to use resources up to a maximum level of usage or service. Storage cluster 102 may enable tenants to use resources within storage cluster 102 up to this maximum level of usage or service (“maximum QoS”) when there are unused resources available within storage cluster 102. Storage cluster 102 may employ a scheduling algorithm, such as the Excessive Weighted Round Robin algorithm (EWRR) algorithm, for admitting work into storage cluster 102. In some examples, storage cluster 102 may make decisions about scheduling at the entry point(s) of storage cluster 102 (e.g., initiator nodes 110) so it is possible to back pressure each of initiator nodes 110 as quickly as possible. Preferably, the scheduling algorithm used ensures that storage cluster 102 allows more work, up to maximum QoS limits, when resources allocated for other volumes are unused.

(42) In general, storage cluster 102 may enforce certain constraints on the number of read operations and write operations of a given fixed block size performed per unit of time for each volume. These operations may be described or measured in terms of the “IOPs,” as noted above. In some examples, constraints on read and write operations may be specified by parameters (each typically expressed in terms of “IOPs”) that are specified when a volume is created. In some examples, independent constraints are provided for both read and write storage cluster operations in terms of IOPs.

(43) For example, “RG” may be the rate of read operations per second that is guaranteed (R=“Read” and G=“Guaranteed”) for a specified volume, assuming, of course, that there is demand that such operations be performed. Therefore, given that there might be no demand for read

operations associated with a specific volume, the actual rate that is guaranteed is the minimum of RG and the actual dynamic read demand being experienced by that specific volume at a particular time. “RM” may be the rate of read operations that storage cluster **102** will not permit to be exceeded (M=“Maximum”), independent of the demand. “WG” may be the rate of write operations per second that is guaranteed for the specified volume (W=“Write”), again assuming, of course, that there is demand. As with the guaranteed read rate, the rate actually guaranteed is the minimum of WG and the dynamic write demand being experienced by that specific volume at a particular time. “WM” is the rate of write operations per second that storage cluster **102** will not permit to be exceeded, independent of the demand.

(44) Separate and independent constraints for read and write operations, as outlined above, may be appropriate at least because the overall processing effort for a write operation may be considerably higher than for a corresponding read operation. For example, a read operation to a non-durable volume might only consume a small number of processor cycles in one DPU **117** (i.e., the DPU **117** the one containing the non-durable volume). However, a write operation to a compressed durable volume will consume more processor cycles writing data to more than one other DPU **117** (e.g., one DPU **117** associated with a primary node, one associated with a secondary node, and one or more associated with plex nodes that are used to store the data). Further, although it may be possible to specify a blended (or mixed) IOPs rate (rather than specifying separate read and write rates), specifying a blended rate is less complete than specifying independent read and write rates.

(45) Note that the terms “guaranteed” and “maximum” may be more accurate descriptions of the above-described terms than “minimum” and “maximum.” Use of the terms “minimum” and “maximum” together might imply that for the minimum rate, the rate does not drop below the specified minimum value. In some implementations, this is not quite accurate, since when there is no demand on a given volume, the rate of operations performed for that volume might be zero.

(46) The quality of service standard may also be adaptive to accommodate dynamic demand for resources within storage cluster **102** at any given time exceeding the total amount of resources offered by storage cluster **102**. For example, it may be appropriate for storage cluster **102** to be oversubscribed, since oversubscribing resources may lead to a more efficient allocation of resources over the long term. It is therefore possible that if all tenants of storage cluster **102** seek to simultaneously use their guaranteed allocation of resources within storage cluster **102**, the aggregate demand for resources could exceed the total resources available within storage cluster **102**. Storage cluster **102** (or, in some cases controller **130**) may detect this excess demand by monitoring the total QoS delivered by storage cluster **102**. If demand exceeds or is close to exceeding available resources, storage cluster **102** may, in some examples, move one or more volumes **121** within storage cluster **102** or to another location. In other examples, storage cluster **102** may adaptively degrade the QoS provided to each of the tenants sharing storage cluster **102**. In most cases, it is advisable to apply such degradation to all tenants within storage cluster **102** in the same way so that each is affected to the same extent.

(47) Applying quality of service standards to storage scenarios has been traditionally performed, if at all, to prioritize storage in storage area networks. However, applying quality of service standards across nodes, DPUs, resources, and/or operation types (e.g., reads, writes, encryption operations, data compression operations, erasure coding operations, other operations) within a storage cluster, particularly one that serving as a scale out and disaggregated storage cluster, as described herein, is particularly complex, but can be performed effectively using the techniques described herein. In particular, techniques described herein enable predictable and adaptive quality of service standards to be achieved effectively in a large scale out disaggregated storage cluster. In addition, techniques described herein may apply to a variety of storage solutions, including but not limited to, block storage, object storage, and file storage.

(48) In FIG. **1B**, and in accordance with one or more aspects of the present disclosure, storage cluster **102** (or data center **101** of FIG. **1A**) may establish quality of service standards for

customers, tenants, operations, and/or resources. For instance, with reference to FIG. 1B, storage cluster **102** may, for a specific tenant, establish a quality of service standard based on a service level agreement associated with or executed by the tenant. In some examples, a quality of service standard may specify, for each tenant, for each storage cluster resource, and/or for each type of operation associated with a given resource, a set of standards that outline performance, availability, capacity, or other expectations associated with services provided by storage cluster **102**. As described herein, the set of standards may specify a guaranteed allocation of performance, availability, capacity, or other metric or attribute for a given resource within storage cluster **102**. Further, the set of standards may specify a maximum allocation of performance, availability, capacity, or other metric or attribute for the given resource.

(49) In some examples, controller **130** may receive information describing the quality of service standards, where the information is from or derived from input originating from an administrator (e.g., through administrator device **133**). In other examples, such input may originate from a representative of the tenant (e.g., through a client device, not specifically shown in FIG. 1B), where the representative selects or specifies attributes of the desired service level agreement or quality of service standard. Quality of service standards may be established for other tenants in the same or a similar way, thereby enabling tenants to customize services provided by storage cluster **102** pursuant to their own needs. In other examples, storage cluster **102** may offer other the same quality of service to each tenant of storage cluster **102**.

(50) Controller **130** may receive a request to allocate a volume. For instance, in an example that can be described with reference to FIG. 1B, controller **130** detects input that it determines corresponds to request to create a new volume. In some examples, the input originates from one or more of initiator nodes **110**, seeking to allocate new storage for a tenant of storage cluster **102** (e.g., tenant “J” or tenant “K” depicted in FIG. 1B). In other examples, the input may originate from an administrator device (e.g., administrator device **133**), which may be operated by an administrator seeking to allocate new storage on behalf of a tenant of storage cluster **102**. In still other examples, the input may originate from a different device.

(51) Controller **130** may allocate a volume. For instance, again referring to FIG. 1B, controller **130** outputs information about the request to allocate a new volume to storage services module **131**. Storage services module **131** evaluates the information and determines that the request is for a new volume is to be allocated for a specific tenant (i.e., tenant “J” in the example being described). Storage services module **131** further determines, based on the input received by controller **130**, information about the volume type and the quality of service to be associated with the new volume. Storage services module **131** accesses data store **132** and determines which of storage nodes **120** may be allocated to supporting the new volume. In some examples, such a determination may involve evaluating which DPUs **117** and storage devices **127** within storage nodes **120** are available to be involved in serving read and write requests to the new volume. In the example being described, storage services module **131** determines that new volume **121J** is to be allocated in response to the input, and further, that volume **121J** is a durable volume allocated using multiple storage nodes **120**. Specifically, storage services module **131** determines that volume **121J** is to be allocated using resources from storage nodes **120A**, **120B**, and **120D** as illustrated by dotted lines radiating from volume **121J** in FIG. 1B. Storage services module **131** causes controller **130** to allocate volume **121J** within storage cluster **102**.

(52) Controller **130** may allocate specific storage nodes **120** as primary, secondary, and plex nodes for volume **121J**. For instance, in some examples, storage services module **131** may allocate volumes using a primary target node along with a number of “plex” nodes. In some implementations, the primary target node may be used as a communication or interaction point for other devices within storage cluster **102** (e.g., initiator nodes **110**). The “plex” nodes may be used to store the data associated with a volume and may be managed by the primary target node. Generally, durable volumes tend to be stored across multiple storage nodes **120** and/or multiple

fault domains and are often stored using an erasure coding or other data durability scheme. In the example being described, storage services module **131** may designate storage node **120A** as a “primary” target node that serves as a primary target or interaction node for operations involving the volume, with one or more of storage nodes **120A**, **120B**, and **120D** serving as “plex” nodes that are used to store data associated with the volume.

(53) Storage services module **131** may also designate one of the three storage nodes **120**, such as storage node **120B**, as a “secondary” target node which serves as a failover target or interaction node when primary storage node **120A** is not available or is inoperable. Typically, storage services module **131** ensures that the designated primary target node (i.e., storage node **120A**) and the secondary target node (i.e., **120B**) are assigned to different storage nodes **120**. Similarly, each of the plex nodes are often stored across different storage nodes **120** or fault domains, but in some cases, the same storage node **120** may be used for both a plex node and the primary target (or, alternatively, as a plex node and the secondary target node). In the example being described, volume **121J** is allocated using three storage nodes **120**, but in other examples, volume **121J** could be stored using a different number of storage nodes **120**.

(54) Controller **130** may update resource information based on the allocation of volume **121J**. For instance, continuing with the example being described with reference to FIG. **1B**, storage services module **131** maintains a counter of available resources within storage cluster **102**, and when allocating volume **121J**, storage services module **131** decrements the counter associated with the resources used for volume **121J**. Storage services module **131** may decrement the counter by an amount corresponding to the functionality need for the type of volume being allocated. In some examples, different numbers of or types of resources may be used for different types of volumes. If a volume is a durable volume, for example, more than one storage node **120** is often used to ensure data redundancy. If a volume is a compressed and/or encrypted volume, other resources such as compression or cryptographic hardware or accelerators within one or more of DPUs **117** may be required. Storage services module **131** may update the appropriate counters in an appropriate way based on the size, type, or other attributes of newly allocated volume **121J**. Storage services module **131** may store the counters or information about the counters within data store **132**. Storage services module **131** may also store, within data store **132**, information about quality of service standards and/or service level agreements associated with volume **121J**. In some examples, such information may involve guaranteed service levels for storage space and processor cycles as well as maximum service levels for storage space and processor cycles.

(55) Controller **130** may allocate subsequent volumes based on availability. For instance, continuing with the example being described in connection with FIG. **1B**, controller **130** detects input that controller **130** determines corresponds to another request to create a new volume. Controller **130** outputs information about the input to storage services module **131**. Storage services module **131** evaluates the input and determines, based on the input, that the new volume is to be allocated for a different tenant (i.e., “tenant K”). Storage services module **131** further evaluates the input and attempts to find a best fit for the new volume (e.g., volume **121K**) based on criteria that includes the number of processing cycles available within each of storage nodes **120** and the storage space available within each of storage nodes **120**. In some examples, storage services module **131** may also seek to allocate resources to the volume **121K** to ensure that the resources associated with the new volume are allocated across multiple fault zones, thereby helping to ensure that any failures associated with resources allocated to the volume have a limited impact on the volume. Similarly, storage services module **131** may also consider other issues, including wear level profiles associated with storage devices **127** within the storage nodes **120** allocated to the new volume, in an attempt to increase the reliability, performance, and/or efficiency of the newly allocated volume. Based on such considerations, storage services module **131** allocates volume **121K** using storage nodes **120B**, **120C**, and storage node **120N** in the example illustrated in FIG. **1B**. Each time controller **130** allocates new volumes, it may update appropriate counters in a

manner that corresponds to the functionality needed for the newly allocated volume (durable, compressed, encrypted, or other attributes). When a volume is deallocated or deleted, controller **130** may return the resources (e.g., storage capacity, processor cycles, SSD cycles, and other resources) associated with each DPU **117** to the pool of resources available to be allocated to new volumes.

(56) In the example described, volume **121J** and volume **121K** are associated with different tenants within storage cluster **102**. Yet storage services module **131** of controller **130** may allocate volume **121K** and volume **121J** such that volume **121K** and volume **121K** share resources associated with one or more of storage nodes **120**. Specifically, in the example being described, both volume **121J** and volume **121K** share resources of storage node **120B** (i.e., DPU **118B** and/or storage devices **127B**), since both volume **121K** and volume **121J** are supported by storage node **120B**.

Accordingly, resources of each of storage nodes **120** may be allocated to multiple volumes **121** across multiple tenants.

(57) Once allocated, storage cluster **102** may schedule an operation to be performed on a volume. For instance, continuing with the example being described in connection with FIG. **1B**, DPU **118A** of initiator node **110A** detects input from server **112A**. DPU **118A** determines that the input corresponds to a request to perform an operation on volume **121J**. In some examples, DPU **118A** may receive the input from server **112A** of initiator node **110A** in response to server **112A** performing processing on behalf of tenant J and seeking to store data (or perform another operation) within storage cluster **102**. In other examples, DPU **118A** may receive the input from server **112A** in response to server **112A** receiving input from an authorized user of tenant J. In either case, DPU **118A** may output a signal over switch fabric **114**. Storage node **120A**, acting as the primary target node for volume **121J**, detects a signal and outputs information about the signal to DPU **117A**. In some examples, if DPU **117A** and/or storage node **120A** are idle, DPU **117A** may cause the operation to be performed. However, in other examples, such as if DPU **117A** and/or storage node **120A** are not idle, DPU **117A** queues the operation (e.g., in a queue maintained by DPU **117A**) to be performed at a later time.

(58) Storage cluster **102** may cause the queued operation to be performed on the volume when resources are available. For instance, continuing with the example being described in connection with FIG. **1B**, DPU **117A** maintains a count (e.g., a credit count) that is used to determine when to dequeue operations from the queue and cause them to be executed. In some examples, the credit count is increased periodically over time, and when sufficient credits are available to perform an operation, an operation is dequeued from a queue and executed. The credit count may be decreased as operations are dequeued from the queue and executed.

(59) As described herein, other storage nodes **120** may serve as a secondary and/or plex node associated with volume **121J**. For instance, when an operation involving volume **121J** is performed, other storage nodes **120** associated with volume **121J** may also have to perform operations. As an example, if storage node **120B** serves as the secondary target node and/or a plex node for volume **121J**, storage node **120B** may have work to do to complete the requested operation. In such an example, DPU **117B** of storage node **120B** may queue operations to be performed at a later time, in a manner similar to that described in connection with DPU **117A** and storage node **120A**. DPU **117B** may periodically increase its own credit count over time, thereby making credits available, and thereby enabling operations to be dequeued and executed. As work is dequeued and executed, DPU **117B** may correspondingly decrease a credit count associated with the resources of DPU **117B** (or storage node **120B**) being used to perform the operation. In this way, operations associated with volume **121J** are eventually performed, in a rate-limited way, by each of storage nodes **120** that are used to implement volume **121J**.

(60) For storage cluster **102**, two primary operations to be performed are write operations and read operations. Both write and read operations are generally associated with a given volume **121**. In addition, both write and read operations for a given volume **121** may involve use of storage

resources across multiple storage nodes **120**.

(61) For example, storage cluster **102** may service a write request associated with an allocated volume. For instance, in another example that can be described in connection with FIG. **1B**, DPU **118A** of initiator node **110A** detects input from server **112A**. DPU **118A** determines that the input corresponds to a request to initiate a write operation to volume **121J**. DPU **118A** outputs a signal over switch fabric **114**. Storage node **120A**, acting as the primary target node for volume **121J**, detects a signal and outputs information about the signal to DPU **117A**. DPU **117A** determines that the signal corresponds to a request to write data to volume **121J**. DPU **117A** further determines that the signal includes data to be written to volume **121J**. DPU **117A** interacts with the plex nodes associated with volume **121J** (e.g., one or more of storage nodes **120A**, **120B**, or **120D**) and writes data to one or more storage devices **127** within storage nodes **120A**, **120B**, and/or storage node **120D**. DPU **117A** may write data to such storage nodes **120** using the data durability attributes, compression attributes, encryption attributes, or other attributes associated with volume **121J**. Each of storage nodes **120** involved in such a write operation may queue work associated with the write operation to be performed pursuant to the quality of service standards (or service level agreements) associated with volume **121J**. If too much work is queued and/or insufficient processing credits are available within one or more of the storage nodes **120** associated with volume **121J**, each such storage node **120** may refrain from processing the work until the queues have less work and/or until sufficient credits are available to enable processing of the work. DPU **117A** also logs information about the write operation to storage node **120B** which serves as a secondary target node for volume **121J**. By writing such information to storage node **120B**, storage node **120B** may be capable of serving as a failover target node should storage node **120A** become unavailable or inoperable. Logging information about the write operation to storage node **120B** may also be performed pursuant to the quality of service standards (or service level agreements) associated with volume **121J**.

(62) Storage cluster **102** may service a read request associated with an allocated volume. For instance, continuing with the example being described with reference to FIG. **1B**, DPU **118B** of initiator node **110A** may detect input from server **112B**. DPU **118B** determines that the input corresponds to a request to read data from volume **121J**. DPU **118B** outputs a signal over switch fabric **114**. Storage node **120A**, acting as the primary target node for volume **121J**, detects a signal and outputs information about the signal to DPU **117A**. DPU **117A** determines that the signal corresponds to a request to read data from volume **121J**. DPU **117A** interacts with one or more of the plex nodes associated with volume **121J** and accesses the requested data from volume **121J**. When DPU **117A** interacts with each of storage nodes **120** involved in the read operation, each involved storage node **120** may queue work associated with the read operation pursuant to the quality of service standards associated with volume **121J** and ensure the read operation is performed in a manner consistent with such standards. After accessing one or more storage nodes **120** associated with volume **121J**, DPU **117A** outputs a signal over switch fabric **114**. DPU **118B** of initiator node **110B** detects a signal over switch fabric **114** and determines that the signal includes the information requested in the original request to read data.

(63) Storage cluster **102** may detect a risk that the quality of service guarantees associated with one or more volumes **121** might not be satisfied. For instance, in some examples, one or more of storage nodes **120** may detect that one or more of volumes **121** allocated and utilized within storage cluster **102** are being utilized to an extent that nears the maximum storage, processing, and/or other capability of storage cluster **102**. Such a situation can arise if controller **130** and/or storage nodes **120** enable oversubscription, which may involve guaranteeing a quality of service standard that storage cluster **102** cannot be simultaneously provide to each of the allocated volumes **121** within storage cluster **102**. Oversubscription may be enabled, however, because in many cases, it is rare that all volumes **121** would require the guaranteed QoS at the same time, and oversubscription enables storage cluster **102** to be used more efficiently over the long run. In some examples,

therefore, controller **130** and/or storage nodes **120** may allow oversubscription of the quality of service standards and/or service level agreements associated with volumes **121**. If oversubscription is enabled, resource conflicts may arise when the number of active volumes **121** collectively require more performance than can be provided by storage cluster **102**. One or more of storage nodes **120** (or controller **130**) may detect a risk of conflicts when the throughput or utilization of one or more resources associated with volumes **121** (e.g., storage space, processing cycles, bandwidth, input/output cycles associated with SSDs or other storage devices) exceeds a threshold. In some cases, the threshold may be just below 100%, such as at 95%.

(64) Storage cluster **102** may move one or more volumes to address the risk that quality of service guarantees might not be satisfied. For instance, in an example that can be described in the context of FIG. **1B**, controller **130** monitors operation of storage cluster **102**. Controller **130** identifies one or more volumes **121** that, if moved, may alleviate the overly-high utilization experienced by the collection of volumes **121** within storage cluster **102**. In one example, controller **130** may identify volume **121K** as a candidate volume that should be moved to address the oversubscription. In other examples, controller **130** may identify additional volumes **121** to move. Controller **130** identifies a destination location for volume **121K**. Such a destination may be within storage cluster **102**, but may involve a different set of storage nodes **120**. In other examples, such a destination may be outside storage cluster **102**. When controller **130** moves volume **121K** to the identified destination, controller **130** may initiate a failover of the volumes so that there is little or no disruption to service within storage cluster **102**. Controller **130** readjusts, based on the move, information about the resources allocated to the affected volumes **121**. Controller **130** updates information about the allocated resources in data store **132**.

(65) Storage cluster **102** may perform an adaptive degradation process in response to determining that quality of service guarantees might not be satisfied. For instance, in an example where controller **130** determines that there is no suitable destination for moving volume **121K**, controller **130** may degrade the quality of service of all volumes **121** adaptively. Often, controller **130** will not know the total demand in such a situation, so controller **130** starts degrading the quality of service standards with a small degradation. If controller **130** detects continued high utilization, controller **130** may continue degrading the quality of service for all of the volumes in small increments until the total throughput across all volumes **121** within storage cluster **102** is below the full capacity of storage cluster **102** minus a hysteresis value. Although such degradation is described as being performed by controller **130**, such a process may be performed by other systems within storage cluster **102**, including storage nodes **120**.

(66) The techniques described herein may provide certain technical advantages. For instance, in examples where a storage cluster is implemented in a data center that is characterized by a scale-out and hyper disaggregated architecture, applying a hierarchical and granular approach to providing quality of service, as described herein, can result in effective use of available resources. In some cases, such effective use of resources may involve associating an allocation of a storage resource with a given tenant, where that allocation is specified in terms of a guaranteed allocation and a maximum allocation of the storage resource. By using resources effectively, tenants, customers, and/or users of the storage cluster experience high throughput, less capacity constraints, and more reliable compliance with quality of service standards and/or service level agreements.

(67) Further, by offloading data processing from general purpose CPUs in a data center to specialized processors (e.g., DPUs **117**), the limits of scaling up the frequency and number of cores in such general purpose CPUs can be avoided or alleviated. Use of DPUs **117** to execute data-centric computations inside initiator nodes **110**, storage nodes **120**, and other systems may result in efficient and/or optimal use of available network resources, and may result in high throughput, less constraints, and more reliable network operation.

(68) Modules illustrated in FIG. **1B** (e.g., storage services module **131**) and/or illustrated or described elsewhere in this disclosure may perform operations described using software, hardware,

firmware, or a mixture of hardware, software, and firmware residing in and/or executing at one or more computing devices. For example, a computing device may execute one or more of such modules with multiple processors or multiple devices. A computing device may execute one or more of such modules as a virtual machine executing on underlying hardware. One or more of such modules may execute as one or more services of an operating system or computing platform. One or more of such modules may execute as one or more executable programs at an application layer of a computing platform. In other examples, functionality provided by a module could be implemented by a dedicated hardware device.

(69) Although certain modules, data stores, components, programs, executables, data items, functional units, and/or other items included within one or more storage devices may be illustrated separately, one or more of such items could be combined and operate as a single module, component, program, executable, data item, or functional unit. For example, one or more modules or data stores may be combined or partially combined so that they operate or provide functionality as a single module. Further, one or more modules may interact with and/or operate in conjunction with one another so that, for example, one module acts as a service or an extension of another module. Also, each module, data store, component, program, executable, data item, functional unit, or other item illustrated within a storage device may include multiple components, sub-components, modules, sub-modules, data stores, and/or other components or modules or data stores not illustrated.

(70) Further, each module, data store, component, program, executable, data item, functional unit, or other item illustrated within a storage device may be implemented in various ways. For example, each module, data store, component, program, executable, data item, functional unit, or other item illustrated within a storage device may be implemented as a downloadable or pre-installed application or “app.” In other examples, each module, data store, component, program, executable, data item, functional unit, or other item illustrated within a storage device may be implemented as part of an operating system executed on a computing device.

(71) FIG. 2A is a block diagram illustrating an example node or data processing unit (DPU) **217**, in accordance with the techniques of this disclosure. DPU **217** of FIG. 2 may generally represent any of nodes **17** of FIG. 1A or DPUs **117** of FIG. 1B. In some examples, DPU **217** represents a hardware chip implemented in digital logic circuitry and may be used in any computing or network device.

(72) If deployed in the context of FIG. 1A and FIG. 1B, DPU **217** may be communicatively coupled to one or more network devices, server devices (e.g., servers **12** or servers **112**), random access memory, storage media (e.g., solid state drives or SSDs), storage devices **127**, a data center fabric (e.g., switch fabric **114**), or the like, e.g., via PCI-e, Ethernet (wired or wireless), or other such communication media. Moreover, DPU **217** may be implemented as one or more application-specific integrated circuit (ASIC), may be configurable to operate as a component of a network appliance or may be integrated with other DPUs within a device.

(73) In the illustrated example of FIG. 2A, DPU **217** includes one or more multi-core processors **135** having a plurality of programmable processing cores **140A-140N** (“cores **140**”) coupled to an on-chip memory unit **134** and one or more storage devices **127A** through **127N** (“storage devices **127**” and representing any number of storage devices). Each of cores **140** may have a level 1 cache **141** (level 1 caches **141A**, **141B**, and **141N** are associated with cores **140A**, **140B**, and **140N**, respectively).

(74) Memory unit **134** may include two types of memory or memory devices, namely coherent cache memory **136**, non-coherent buffer memory **138**, and non-volatile memory **139** (e.g., NVDIMM memory). Processor **135** also includes a networking unit **142**, work unit (WU) queues **143**, a memory controller **144**, and accelerators **146**. Although not shown, processor **135** may also include a storage device controller used when accessing storage devices **127**. As illustrated in FIG. 2A, each of cores **140**, networking unit **142**, WU queues **143**, memory controller **144**, memory unit

134, and accelerators **146** are communicatively coupled to each other. In some examples, processor **135** of DPU **217** further includes one or more accelerators configured to perform acceleration for various data-processing functions, such as look-ups, matrix multiplication, cryptography, compression, regular expressions, or the like.

(75) In the example of FIG. 2A, DPU **217** represents a high performance, hyper-converged network, storage, and data processor and input/output hub. For example, networking unit **142** may be configured to receive one or more data packets from and transmit one or more data packets to one or more external devices, e.g., network devices. Networking unit **142** may perform network interface card functionality, packet switching, and the like, and may use large forwarding tables and offer programmability. Networking unit **142** may expose Ethernet ports for connectivity to a network, such as switch fabric **114** of FIG. 1A or FIG. 1B. DPU **217** may also include one or more interfaces for connectivity to host devices (e.g., servers) and data storage devices, e.g., solid state drives (SSDs) via PCIe lanes. DPU **217** may further include one or more high bandwidth interfaces for connectivity to off-chip external memory.

(76) Processor **135** further includes accelerators **146** configured to perform acceleration for various data-processing functions, such as look-ups, matrix multiplication, cryptography, compression, data durability and/or reliability, regular expressions, or the like. For example, accelerators **146** may comprise hardware implementations of look-up engines, matrix multipliers, cryptographic engines, compression engines, or the like. In accordance with the some implementations described herein, accelerators **146** may include a data durability unit that may be used to implement one or more data durability and/or reliability schemes. In some examples, such a data durability unit may be configured to perform matrix multiplication operations commonly performed in erasure coding schemes, such as Reed Solomon erasure coding schemes. Such a data durability unit may be configured to efficiently perform operations, such as those relating to Galois Field mathematics, that might be difficult and/or inefficient to perform using commonly available processors or other processing hardware. Further, such a data durability unit may be designed to perform and/or implement multiple different types of data durability schemes by configuring different matrices specific to each implementation.

(77) Memory controller **144** may control access to on-chip memory unit **134** by cores **140**, networking unit **142**, and any number of external devices, e.g., network devices, servers, external storage devices, or the like. Memory controller **144** may be configured to perform a number of operations to perform memory management in accordance with the present disclosure. For example, memory controller **144** may be capable of mapping accesses from one of the cores **140** to either of coherent cache memory **136** or non-coherent buffer memory **138**. More details on a bifurcated memory system that may be included in DPU **217** are available in U.S. Pat. No. 10,565,112, filed Apr. 10, 2018, and titled "Relay Consistent Memory Management in a Multiple Processor System," the entire content of which is incorporated herein by reference.

(78) Cores **140** may comprise one or more microprocessors without interlocked pipeline stages (MIPS) cores, advanced reduced instruction set computing (RISC) machine (ARM) cores, performance optimization with enhanced RISC-performance computing (PowerPC) cores, RISC Five (RISC-V) cores, or complex instruction set computing (CISC or x86) cores. Each of cores **140** may be programmed to process one or more events or activities related to a given data packet such as, for example, a networking packet or a storage packet. Each of cores **140** may be programmable using a high-level programming language, e.g., C, C++, or the like.

(79) In FIG. 2A, and in accordance with one or more aspects of the present disclosure, data processing unit **217** may process an operation in a storage cluster. For instance, networking unit **142** of data processing unit **217** receives an indication of work associated with a read or write operation associated with a data storage unit (e.g., a "volume") in a storage cluster, such as storage cluster **102** of FIG. 1B. Networking unit **142** outputs information about the work to one or more of cores **140**. One or more of such cores **140** queues the work to be processed by data processing unit

217 and/or by another data processing unit 217 within the storage cluster. In some examples, data processing unit 217 ensures the work is performed pursuant to quality of service standards associated with the relevant volume specified by the work. Data processing unit 217 may queue the work for later processing using a mechanism that ensures that quality of service standards (e.g., the guaranteed IOPs associated with the relevant volume) are met. If one or more queues associated with the relevant volume do not hold any other work, operations associated with the work may be performed without queueing the work for later processing. In such an example, data processing unit 217 may ensure that the work is not performed at a rate that exceeds the maximum quality of service level allowed by the QoS standards associated with the volume.

(80) FIG. 2B is a conceptual diagram illustrating one possible partitioning of virtual processors within an example processor, in accordance with one or more aspects of the present disclosure. FIG. 2B illustrates processor 135, which may correspond to processor 135 included within data processing unit 217 of FIG. 2A. As illustrated in FIG. 2B, processor 135 includes an array of processors corresponding to those included within data processing unit 217 of FIG. 2A. The processors illustrated in FIG. 2B are shown as an array of cores 140. Although the processors are illustrated as cores 140, such processors may alternatively be physical processors, virtual processors, hardware threads included within one or processing cores, or other processing units.

(81) To support the significant bandwidth preferred for both frontend and backend operations (e.g., 400 Gbps bandwidth or more), the available processing cycles associated with the cores may be partitioned into 3 pools: networking pool 202, storage or volume pool 204, and SSD or backend pool 206. In the example of FIG. 2B, 48 cores 140 are illustrated, and each is allocated to one of these three functions. Each of the allocated functions may also correspond to a level in a hierarchy of resources, where networking pool 202 corresponds to NVMe/TCP operations, volume pool 204 corresponds to operations performed by the volume layers, and backend pool 206 corresponds to operations performed by the storage devices (e.g., SSDs) and/or the backend layer. SSD pool 208 is also included within core hierarchy 210 and represents the SSD resources included within a storage cluster.

(82) Processing performed by the volume cores 140 (e.g., volume pools 204) tends to be the source of any bottleneck, so it may be appropriate to help ensure that the computational resources allocated to the cores used to service volume-related operations are sufficient, such as by allocating more resources to volume pool 204. Therefore, in the specific example illustrated in FIG. 2B, 16 cores are allocated to networking pool 202, 24 cores are allocated to volume pools 204, and 8 cores are allocated to backend pools 206. In addition, the scheduling logic (e.g., logic for scheduling operations across storage nodes 120 of FIG. 1B) may be implemented on some of cores 140 from networking pools 202. Partitioning computational resources as illustrated in FIG. 2B may enable quality of service standards to be supported for many volumes (e.g., on the order of thousands or more) on a per-volume basis.

(83) Accordingly, within each data processing unit 217, the CPU cores may be hard partitioned to perform separate processing operations (i.e., network, volume, and back-end). This partitioning may be done in a way that ensures that cores 140 serving the network and backend functions will not likely become a bottleneck for any of the scenarios considered, even if less computational resources are allocated to networking pool 202 and backend pool 206.

(84) There may be a wide variation in performance expectations for a given volume (e.g., on the order of 1K to 2M IOPs). As described below in connection with FIG. 5A and FIG. 5B, managing quality of service standards may be performed using a hierarchical tree of resources with a set of queues acting as leaf nodes in the tree. To support high performance volumes from the perspective of one or more of initiator nodes 110 (see FIG. 1B), multiple queues per volume may need to be supported. In some examples, this may mean translate into one TCP connection per queue. In some cases, however, multiple queues can be merged into a single queue (e.g., when using FCP protocol, described in United States Patent Publication No. 2019-0104206 A1, entitled "FABRIC CONTROL

PROTOCOL FOR DATA CENTER NETWORKS WITH PACKET SPRAYING OVER MULTIPLE ALTERNATE DATA PATH,” and which is hereby incorporated by reference).

(85) In some examples, code or logic that manages and/or ensure compliance with QoS standards may execute on some of the cores **140** assigned to volume pool **204**. For instance, if 32K queues (see, e.g., FIG. 5A and FIG. 5B) are supported, and 32 processors/cores/virtual processors are allocated to networking pool **202**, each processor/core/virtual processor may be responsible for 1000 queues. However, as QoS are managed on a per-volume basis, all the queues belonging to a volume may be managed together, which may require that all of those queues should be handled by a single processing unit to be performance and/or cache efficient. When a single volume requires a very high number of IOPs, the queues belonging to that volume may be split into a few groups as a single networking processing unit (e.g., core **140**) might not have enough processing capability. In such an example, the required QoS may be split among the groups with each group managed by its own QoS rate limiter/shaper. In some examples, the QoS rate limiter/shaper may execute the EWRR algorithm, referenced above. When the QoS of a volume is split among multiple shapers, all these shapers together will typically provide the required QoS by sharing the available credits. All the shapers belonging to a volume (or a shared volume which is a volume shared with multiple initiators) will typically consult a common process periodically for sharing the credits.

(86) Volume aggregates may be needed to support QoS for snapshots and clones and also for high performance volumes that may be implemented using striping. It may be possible to handle volume aggregates within a three level QoS hierarchy (see FIG. 5A and FIG. 5B) by having the storage service operations (e.g., those handled by controller **130**) statically distribute the guaranteed and maximum (G, M) values for the aggregate over the volumes that are part of the aggregate. Therefore, even though there may be four levels of abstraction (queues, volumes, volume aggregates, and root), it may be possible to handle volume aggregates with only three levels in the QoS hierarchy.

(87) FIG. 3 is a block diagram illustrating another example of an example DPU, in accordance with one or more aspects of the present disclosure. DPU **317** of FIG. 3 may operate substantially similar to and generally represent any of nodes **17** of FIG. 1A or DPUs **117** of FIG. 1B. Therefore, like data processing unit **217** of FIG. 2A, if DPU **317** is deployed in the context of FIG. 1A and FIG. 1B, data processing unit **317** may be communicatively coupled to a data center fabric (e.g., switch fabric **114**), one or more server devices (e.g., servers **12** or servers **112**), storage media (e.g., SSDs), storage devices **127**, one or more network devices, random access memory, or the like, e.g., via PCI-e, Ethernet (wired or wireless), or other such communication media in order to interconnect each of these various elements. In some examples, DPU **317** may represent a hardware chip implemented in digital logic circuitry. DPU **317** may also be implemented as an integrated circuit mounted on a motherboard of a computing, networking and/or storage device or installed on a card connected to the motherboard of the device.

(88) As with other implementations of example DPUs described and/or illustrated herein, DPU **317** may represent a high performance, hyper-converged network, storage, and data processor and input/output hub. As illustrated in FIG. 3, DPU **317** includes networking unit **152**, processing clusters **156A-1** to **156N-M** (processing clusters **156**), host units **154A-1** to **154B-M** (host units **154**), and central cluster **158**, and is coupled to external memory **170**. Each of host units **154**, processing clusters **156**, central cluster **158**, and networking unit **152** may include a plurality of processing cores, e.g., MIPS cores, ARM cores, PowerPC cores, RISC-V cores, or CISC or x86 cores. External memory **170** may comprise random access memory (RAM), dynamic random access memory (DRAM), or non-volatile memory.

(89) As shown in FIG. 3, host units **154**, processing clusters **156**, central cluster **158**, networking unit **152**, and external memory **170** are communicatively interconnected via one or more specialized network-on-chip fabrics. A set of direct links **162** (represented as dashed lines in FIG. 3) forms a signaling network fabric that directly connects central cluster **158** to each of the other

components of DPU **317**, that is, host units **154**, processing clusters **156**, networking unit **152**, and external memory **170**. A set of grid links **160** (represented as solid lines in FIG. 3) forms a data network fabric that connects neighboring components (including host units **154**, processing clusters **156**, networking unit **152**, and external memory **170**) to each other in a two-dimensional grid. (90) Networking unit **152** has Ethernet interfaces **164** to connect to the switch fabric, and interfaces to the data network formed by grid links **160** and the signaling network formed by direct links **162**. Networking unit **152** provides a Layer 3 (i.e., OSI networking model Layer 3) switch forwarding path, as well as network interface card (NIC) assistance. One or more hardware direct memory access (DMA) engine instances (not shown) may be attached to the data network ports of networking unit **152**, which are coupled to respective grid links **160**. The DMA engines of networking unit **152** are configured to fetch packet data for transmission. The packet data may be in on-chip or off-chip buffer memory (e.g., within buffer memory of one of processing clusters **156** or external memory **170**), or in host memory.

(91) Host units **154** each have PCI-e interfaces **166** to connect to servers and/or storage devices, such as SSD devices. This allows DPU **317** to operate as an endpoint or as a root. For example, DPU **317** may connect to a host system (e.g., a server) as an endpoint device, and DPU **317** may connect as a root to endpoint devices (e.g., SSD devices). Each of host units **154** may also include a respective hardware DMA engine (not shown). Each DMA engine is configured to fetch data and buffer descriptors from host memory, and to deliver data and completions to host memory.

(92) DPU **317** may provide optimizations for stream processing. For instance, DPU **317** may execute an operating system that facilitates run-to-completion processing, which may eliminate interrupts, thread scheduling, cache thrashing, and associated costs. For example, an operating system may run on one or more of processing clusters **156**. Central cluster **158** may be configured differently from processing clusters **156**, which may be referred to as stream processing clusters. In one example, central cluster **158** executes the operating system kernel (e.g., Linux kernel) as a control plane. Processing clusters **156** may function in run-to-completion thread mode of a data plane software stack of the operating system. That is, processing clusters **156** may operate in a tight loop fed by work unit queues associated with each processing core in a cooperative multi-tasking fashion.

(93) FIG. 4 is a conceptual diagram illustrating various roles played by one or more nodes in a storage cluster, in accordance with one or more aspects of the present disclosure. FIG. 4 can be described in the context of FIG. 1B. For example, in FIG. 1B, storage services module **131** of controller **130** performs various storage cluster services. Controller **130** of FIG. 1B may correspond to cluster services node **430** of FIG. 4. Also, in the context of FIG. 1B, initiator node **410** corresponds to one of initiator nodes **110**, while primary node **421**, secondary node **422**, and plex nodes **423** correspond to one or more of storage nodes **120**. As illustrated in FIG. 1B, the overall functionality of storage cluster **102** may be divided into control operations that are performed by cluster services node **430** (i.e., controller **130** of FIG. 1B), and data operations that are executed by initiator node **410** (i.e., initiator nodes **110** of FIG. 1B) and primary node **421**, secondary node **422**, and plex nodes **423** (i.e., storage nodes **120** of FIG. 1B). The cluster service node **430** communicates with all DPU based nodes. In this case, it may also communicate with some or all of the primary, plex, and secondary nodes.

(94) The control operations performed by cluster services node **430** in FIG. 4 are primarily operations to create and delete volumes (i.e., the basic unit of storage abstraction). Cluster services node **430** may, however, also perform other operations. In general, control operations may be performed by cluster services node **430** communicating with other DPU-based nodes (e.g., primary, secondary, and plex nodes, along with DPU-based initiator nodes). The control operations may be initiated by users via a higher level orchestration system using an intent based API provided by the cluster services node **430**. In response, the cluster services node **430** configures the storage nodes of the cluster as appropriate. As discussed below, read/write operations are often data operations

initiated by the initiator nodes **110**. The “ack” communications illustrated in FIG. 4 (flowing from cluster services node **430** to initiator **410**) may be considered acknowledgements of those read/write operations. The functions performed by cluster services node **430** when creating a volume may include “allocation,” which is the process of assigning resources (e.g., DPUs **117**, storage devices **127**, bandwidth, and/or other network resources) on behalf of a volume at the time a volume is created by initiator node **410**. Cluster services node **430** performs allocation operations when creating a volume, and cluster services node **430** may also later make adjustments to the initial allocation of resources associated with a volume. Such adjustments may be prompted by, for example, a need to move a volume to a different location (e.g., to a different set of storage nodes **120** within storage cluster **102**). Such adjustments may also be prompted by a volume being deleted, with the adjustments in such a case resulting in resources being released back to the storage cluster. Allocation typically involves multiple resource types distributed across the network and may also involve taking other constraints into account (e.g., fault zones and wear profiles associated with physical storage devices). In some examples, the problem to be solved by allocation corresponds to or is equivalent to a problem sometimes referred to as multidimensional bin packing with constraints. Therefore, techniques and algorithms available for solving problems characterized as multidimensional bin packing with constraints may also be applied to allocation.

(95) Data operations performed by initiator node **410**, primary node **421**, secondary node **422**, and plex nodes **423** include read and write operations. Read and write operations may also be initiated by initiator node **410**. In general, initiator node **410** initiates a read or write operation by interacting with primary node **421**. Should primary node **421** be offline or otherwise unavailable, initiator node **410** interacts with secondary node **422**. A read operation performed in the context of FIG. 4 may involve primary node **421** accessing data stored within at least some of plex nodes **423**. For example, a durable volume may have data stored across multiple plex nodes **423** for redundancy purposes, so primary node **421** may access multiple plex nodes **423** in order to read the requested data. Similarly, a write operation may involve primary node **421** writing data to one or more of plex nodes **423**. For a durable volume that implements a redundancy scheme (e.g., an erasure coding scheme), primary node **421** may write data to multiple plex nodes **423** within a storage cluster.

(96) For write operations for a given volume, primary node **421** may also log information about the write operation to secondary node **422**. If secondary node **422** is called upon to perform operations associated with the volume, secondary node **422** may therefore have the information needed to perform any associated with the volume.

(97) Resources within a storage cluster are not often allocated exclusively to a volume, so they typically are shared dynamically across volumes. Sharing such resources may be performed through a process called scheduling. In some examples, scheduling may involve sharing resources dynamically (i.e. at the time of executing a read or write operation) with the goal of satisfying potentially competing constraints requested by active volumes within a storage cluster. Unlike allocation functions that are performed by initiator node **410**, scheduling may, in some examples, be performed entirely by primary node **421**, secondary node **422**, and/or plex nodes **423**. In other words, in the context of FIG. 1B, scheduling may be performed entirely by storage nodes **120** (or DPUs **117** within storage nodes **120**), where DPUs **117** collaborate to schedule the use of network resources at run time. In some examples, the process of allocation might be considered to be static relative to read and write operations, while the scheduling process might be considered to be dynamic.

(98) In FIG. 4, note that various DPUs **117** (and DPUs **118**) included within corresponding initiator nodes **110** and storage nodes **120** (see FIG. 1B) may play different roles (e.g., initiator node **410**, primary node **421**, secondary node **422**, plex node **423**). In some examples, in order to maintain an appropriate quality of service level, it may be important to segregate the traffic received by primary nodes **421** from initiator nodes **410** (new work) from the traffic generated among DPUs **117** (e.g., corresponding to write operations from primary nodes **421** to secondary nodes **422** and primary

nodes **421** to plex nodes **423**). By segregating such traffic, it is possible to ensure that work admitted as guaranteed by QoS standards does not get backed up anywhere in the cluster by other work not subject to such QoS standards.

(99) The rate at which read and write operations are executed can depend on many types of factors and resources. In practice, however, there is often one resource that turns out to be a bottleneck, and the availability of this resource tends to determine the service rate. In a data flow graph such as that illustrated in FIG. 4, it is unlikely that bottlenecks will develop simultaneously in multiple places during the processing of a given operation. Further, it is likely that different instances of the same operation will encounter the same bottleneck as will instances of other operations.

(100) For example, each of DPUs **117** of FIG. 1B may have many internal resources. Such resources may include customized embedded processors with multiple threads or virtual processors (VPs), hardware accelerators (e.g., capable of performing encryption, compression, and other cryptographic operations) coherent and non-coherent memory operations involving on-chip and DDR memories, DMA engines, SSD processing capability and bandwidth. Yet it may be appropriate to make the simplifying assumption that processor cycles (i.e., VP cycles) tend to limit resource availability within DPUs **117**, at least in so far as storage operations are concerned. Also note that the processor or VP cycles needed for a given read or write operation are independent of the logical block size.

(101) In FIG. 4, when allocating a volume, cluster services node **430** needs to know the amount of resources (e.g., processor cycles) expended per read and write for each of three DPUs **117** for each type of volume. There may be a significant number of volume types (e.g., durable, nondurable, encrypted, nonencrypted, compressed, not compressed, and volume types with any combination of these attributes). Therefore, there may be six static values for each volume type, which can be expressed as a 3×2 work matrix WM.sub.T, where T is the volume type:

(102) TABLE-US-00001

		Read	Write
	Primary(<i>P</i>)	[PR	PW]
WM _T =	Secondary(<i>S</i>)	[SR	SW]
	Plex(<i>X</i>)	[XR	XW]

units: kilo-instructions

(103) In the matrix above, when no VP cycles are expended for an operation (e.g., the S and X column values for a non durable read), the corresponding values may be zero. When a volume is created, cluster services node **430** (e.g., based on input provided by a user of cluster services node **430**) must supply the volume type T and the guaranteed and maximum rates for each of read and write (e.g., in a form such as the matrix above). In some examples, these rates may be termed GRIOPS (Guaranteed Read IOPs), MRIOPS (Maximum Read IOPs), GWIOPS (Guaranteed Write IOPs), and MWIOPS (Maximum Write IOPs). Cluster services node **430** may first translate the read and write guarantees (GRIOPS and GWIOPS) into instruction budgets in each of primary node **421**, secondary node **422**, and plex nodes **423** by using the matrix for the specified volume type. Cluster services node **430** may then use the instruction budgets to allocate processor cycles in one or more DPUs **117** within primary node **421**, secondary node **422**, and plex nodes **423**. Once the allocation is performed, cluster services node **430** may updates the number of processor cycles available in the chosen DPUs **117**. When a volume is deleted, cluster services node **430** returns the resources (i.e., processor cycles) allocated to the volume being deleted to the free pool for each of the affected DPUs **117**.

(104) FIG. 5A is a conceptual diagram illustrating a model for scheduling a single resource within an example storage cluster, in accordance with one or more aspects of the present disclosure. FIG. 5A illustrates a multilevel hierarchical tree of resources terminating at a root node, which may be root node **501** shown at the root of the tree. Root node **501** has a number of child nodes **502**, and each child node **502** includes one or more queues **503**, which themselves may be considered child

nodes of child nodes **502** or leaf nodes of the diagram illustrated in FIG. 5A. For example, queues **503A** are considered child nodes of child node **502A**, queues **503B** are considered child nodes of child node **502B**, and in general, queues **503N** are considered child nodes of child node **502N**. Although the tree illustrated in FIG. 5A is drawn with three levels, any number of levels might be used.

(105) Each of queues **503** represents a data structure in which each queue entry represents a demand on a storage cluster resources. As suggested above, since processor cycles or processor resources (e.g., IOPs) are often a bottleneck in a storage cluster, for convenience in FIG. 5A, each queue entry is assumed to be a demand on processor cycles (IOPs) from processors within volume pool **204** of FIG. 2B.

(106) The total number of processor cycles available is represented by two symbols, Γ and M , which are associated with root node **501**. Γ is the number of processor cycles guaranteed to be available at all times, while M is the maximum number available. For ease in illustration and explanation, Γ is assumed to be equal to M , although in other examples Γ might not be equal to M .

(107) In the example of FIG. 5A, each node in the hierarchy is responsible for allocating its resources, and only its resources, to its children. If root node **501** has (Γ, M) as guaranteed and maximum resources, and child nodes **502** have G_i and M_i as guaranteed and maximum, the resources associated with root node **501** are considered to be oversubscribed if $\sum G_i > \Gamma$ and non-oversubscribed if: $\sum G_i \leq \Gamma$

(108) An oversubscription factor may be defined as: $\sum G_i / \Gamma$.

(109) In FIG. 5A, work arrives at the queues at the top of the diagram, with each entry in a queue representing a specific demand on processor cycles. Each child node is free to ask for as much of the resources as it wishes. This ask may be dynamic and might be represented as “d.sub.i” herein. The child's parent will distribute its guaranteed resources Γ to its children according to an appropriate algorithm, which may be the excess weighted round robin (EWRR) algorithm. Typically, for each level in FIG. 5A, such an algorithm ensures that: 1. It is work conserving: i.e. as long as $\sum d_{sub.i} \geq \Gamma$ no resource units are wasted. 2. For $F \leq 1$, each child gets $\min(d_{sub.i}, G_{sub.i})$. 3. For $\Gamma > 1$, each child gets $\min(d_{sub.i}, G_{sub.i}^*)$ where $G_{sub.i}^* = G_{sub.i} / F$. 4. No child ever gets more resource than its maximum $M_{sub.i}$.

(110) In some examples, QoS control may be provided at the level of volumes only, and not at the granularity of queues. In such an example, the queues for a given volume may be treated identically. This enables storage cluster **102** to use a single scheduler (e.g., executing the EWRR algorithm) operating at the level of volumes.

(111) FIG. 5B illustrates a storage data path distributed across physically distinct nodes or data processing units, in accordance with one or more aspects of the present disclosure. In the example shown, each of primary storage DPU **521**, secondary storage DPU **522**, and plex storage DPU **523** may correspond to one of storage nodes **120** for a given set of volumes. Each of primary storage DPU **521**, secondary storage DPU **522**, and plex storage DPU **523** serves all three functions (primary, secondary, and plex) on behalf of different volumes. Accordingly, the resources of a given one of the DPUs should be properly accounted for and allocated by controller **130** (see FIG. 1B) and properly scheduled by the QoS processing in the DPU. In the example of FIG. 5B, each of primary storage DPU **521**, secondary storage DPU **522**, and plex storage DPU **523** have $P+2$ volume nodes in each DPU, where “P” is the number of volumes for which the DPU acts as the primary DPU (e.g., corresponding to primary storage DPU **521**). The other two volume nodes (“+2”) represent plex and secondary nodes. Note that in the example shown, there is only one plex node and one secondary node per DPU independent of the number of volumes.

(112) FIG. 5B suggests that the secondary and plex volume functions (e.g., secondary storage DPU **522** and plex storage DPU **523**) are deeper in the data path than primary functions. To avoid blockages in the data path, it may be appropriate to provide sufficient resources (e.g., processor cycles) for the secondary and plex volume functions. Otherwise, there is a risk that the primary

volume functions could eventually back up and cause an under-utilization of resources and consequently, poor performance. Providing sufficient resources can be accomplished by slightly over-provisioning the guaranteed resources (e.g., processor cycles) “Gs” and “Gx” for these functions. In some examples, it may be appropriate and/or sufficient to over-provision by a factor of 1.1 to 1.5.

(113) In some examples, an implementation of a scheduler/shaper/rate limiter, such as one that may schedule work pursuant to quality of service standards in FIG. 5B, can be broken into three independent processes: (1) an enqueue process, (2) a dequeue process, and a (3) a credit process. One instance of each of these processes may execute on each of the networking processors (e.g., cores **140** allocated to networking pool **202**) that participate in managing quality of service compliance. If implementations of the process on each of such cores **140** are identical, the implementation can be described from the point of view of a single networking core **140**.

(114) In particular, a data structure may be maintained per volume (where each volume has a quality of service shaper), which may have the form illustrated in the table below:

(115) TABLE-US-00002 Parameter Width Description Static Parameters (in bits) g_cdt 32 Guaranteed credits added every time interval. MAX (Guarantee_Read_iops, Guarantee_Write_iops * SSD R/W) p_cdt 32 Peak credits added every time interval MAX(Peak_Read_iops, Peak_Write_iops * SSD R/W) p_clip 32 Peak credit CLIP (This is to implement burst IOPs) (granularity is IOPs) rd_cdt 32 Credits to consume if the operation is a read operation. wr_cdt 32 Credits to consume if the operation is a write operation. Dynamic Parameters (in bits) active_g_cdt 32 Available guaranteed credit in terms of IOPs active_p_cdt 32 Available peak credit in terms of IOPs TOTAL 224 Total volume bits/bytes (~32B)

(116) A queue, such as one of queues **503** illustrated in FIG. 5A, may be used to hold pending transactions for a given QoS index. Such a queue may be indexed by a quality of service index as follows: pending_q[qos_idx]

(117) In some examples, a 1 bit per QoS index may be used indicate whether or not there are pending transactions with positive active guarantee credits: g_pending

(118) Also, a 1 bit per QoS index may be used to indicate that there are pending transactions with positive active peak credits (“p_pending”). There may also be a 1 bit per QoS index to indicate whether it is free.

(119) When a volume is created, controller **130** of FIG. 1B may create an instance of the QoS shaper using function having a signature such as that illustrated below. Such a function call will update the respective fields in the QoS data structure of the processor and return the QoS index that is allocated for this volume.

(120) TABLE-US-00003 qos_idx_t create_volume_qos(g_cdt, p_cdt, p_clip, rd_cdt, wr_cdt);

(121) The QoS credits may be updated using the below function call whenever there is a need to change the QoS parameters.

(122) TABLE-US-00004 void update_volume_qos(qos_idx, g_cdt, p_cdt, p_clip, rd_cdt, wr_cdt);

(123) A QoS shaper may be deleted using a function call having a signature such as that illustrated below: void delete_volume_qos(qos_idx);

(124) The enqueue process may be implemented such that it is asynchronous to the other two processes (dequeue and credit) and may be triggered when DPU **117** receives a request over the network. The QoS index is determined from the volume index which itself is derived from the queue in which the work arrived. Below is a function to enqueue work to the QoS scheduler; the function outlined below returns true if the volume queues can go forward without stalling, else it returns false and queues the context into the pending queue. In this implementation, the “new_enq” parameter will be true only for enqueues that are not followed by dequeue, to flag whether an enqueue credit check is to be performed.

(125) TABLE-US-00005 bool enqueue(qos_idx, new_enq, context, operation = read/write,

```

num_io); // Determine the credits needed for the enqueue dec_cdt = (operation == read) ? credit
[qos_idx].rd_cdt * num_io : credit[qos_idx].wr_cdt * num_io // If no transactions pending and if
the guarantee credits are positive, we can proceed further if ( (new_enq && (pending_q[qos_idx]
is empty) or (!new_enq && g_pending == 0) ) && credit [qos_idx].active_g_cdt is positive) {
    credit[qos_idx].active_g_cdt -= dec_cdt    credit[qos_idx].active_p_cdt -= dec_cdt
return true } // no pending transactions with positive guarantee credits and it has peak credits, we
can proceed further if ( (new_enq && (g_pending == 0) or (!new_enq && p_pending == 0) )
&& (pending_q[qos_idx] is empty) && credit[qos_idx].active_p_cdt is positive) {
    credit [qos_idx] .active_p_cdt -= dec_cdt    return true } // Push the context into a queue if the
work cannot be admitted into the cluster push (context, operation, num_io) to
pending_q[qos_idx]

```

(126) The Dequeue process may be invoked by a timer, and at each epoch, the process dequeues any pending IOPs for the volumes where credits are available:

```

(127) TABLE-US-00006 // Return if there are no credits if(g_pending == 0 && p_pending == 0) {
    return } // Dequeue if guarantee credits are available if (g_pending != 0) { // find the first
index from g_last_idx that is non_zero that is denoted by g_idx tran = pop the head entry from
pending_q[g_idx] dec_cdt = (tran.operation == read) ? credit[g_idx].rd_cdt * tran.num_io :
credit[g_idx].wr_cdt * num_io    credit[g_idx].active_g_cdt -= dec_cdt
credit[g_idx].active_p_cdt -= dec_cdt    if(credit[g_idx].active_g_cdt is negative ||
(pending_q[g_idx] is empty) )    g_pending[g_idx] = 0    if(credit[g_idx].active_p_cdt is
negative || (pending_q[g_idx] is empty) )    p_pending[g_idx] = 0    g_last_idx = g_idx
deque_callback(tran.context, tran.num_io) } // Dequeue if peak credits are available else if
(p_pending != 0) {    find the first index from p_last_idx that is non zero that is denoted by
p_idx tran = pop the head entry from pending_q[p_idx] dec_cdt = (tran.operation == read) ?
credit[p_idx].rd_cdt * tran.num_io : credit[g_idx].wr_cdt * num_io    credit[p_idx].active_p_cdt
-= dec_cdt    if(credit[p_idx].active_p_cdt is negative || (pending_q[p_idx] is empty) )
p_pending[p_idx] = 0    p_last_idx = p_idx    deque_callback(tran.context, tran.num_io) }

```

(128) The credit process may be triggered by a timer and in general, is responsible for updating the active credits. The interval between timer events should be chosen carefully to satisfy two competing constraints. The first constraint is that the granularity of credit increments be as small as possible to ensure smooth behavior for each queue. This requires the timer interval to be small. The second constraint is that the number of networking VP cycles spent by this process be a small fraction (e.g., <10%) of the total networking VP cycles available.

```

(129) TABLE-US-00007 Initialize update_idx = 0 // check that it does not overflow
credits[update_idx].active_g_cdt += credits[update_idx].g_cdt // clamp the guarantee
credits[update_idx].active_p_cdt += credits[update_idx].p_cdt if( (credits[update_idx].active_p_cdt
overflow) or (credits[update_idx].active_p_cdt > credits[update_idx].p_clip) ) {
    credits[update_idx].active_p_cdt = credits[update_idx].p_clip }
if(credits[update_idx].active_g_cdt is positive && pending_q[qos_idx] is not empty) {
    g_pending[update_idx] = 1 }    if(credits[update_idx].active_p_cdt is positive &&
pending_q[qos_idx] is not empty) {    p_pending[update_idx] = 1 } update_idx++ // if done
with all the qos indexes then wrap back to 0 if(update_idx == last_idx) update_idx = 0

```

(130) In one example of computing the timer interval for the credit process, each processor can be assumed to process 600M instructions per second. One percent of this budget is 6M instructions per second. Assuming 512 volumes per processor, that translated into 12K instructions per volume-second. If each pass takes 100 instructions per volume and 200 instructions for the queues in the volume, the number of passes in 1 second for supporting 12K volumes is 40 (=12000/300). In such an example, the timer interval may be chosen to be 25 milliseconds.

(131) For mixed workloads involving both read and write operations, the QoS that is guaranteed will be somewhere between the read QoS guaranteed and the write QoS guarantee. Calculating the

specific QoS guarantee may be performed by using a ratio of expected reads to writes in the workload, and applying that ratio to the read and write QoS guarantees. In an example using a single shaper per volume, with credits equal to maximum of the read and write QoS guarantees, the credit may be decremented by 1 for every read and the credit may be decremented by $1 \cdot RG/WG$ (ratio of read/write guarantees). Such an example may assume the read QoS guarantee is higher than write QoS guarantee, which is usually the case since the total IOPs available from a typical SSDs is consistent with that assumption. Using such a calculation ensures that the QoS shaper provides RG when the workload is read-only, WG when the workload is write-only, and an allocation of resources in between RG and WG for mixed workloads, where the specific allocation may depend on the ratio of reads to writes in the IO demand. A similar application of this logic can be used for the credits used for enforcing QoS maximum limits.

(132) In some examples, initiator node **110** and storage node **120** may rate limit demand for resources through queue sizing. For instance, in an example that can be described with reference to FIG. 1B, initiator node **110** may discover capabilities of a target storage node **120** when establishing a connection with that storage node **120** using NVMe protocol. Such information may include information about the number of queues and queue depth for storage nodes **120**. A maximum QoS value associated with a given storage node **120** may be included in the capabilities reported by the storage node **120** and this information may be used to achieve coarse level rate limiting at initiator node **110**. In some examples, storage node **120** may take a worst case latency from initiator node **110** into account when reporting the number of queues supported and the queue depth of each of such queues.

(133) Further, in some examples, storage node **120** may also rate limit demand for resources through TCP window size configurations. For instance, again in an example that can be described with reference to FIG. 1B, storage node **120** may use TCP protocol tuning to limit the maximum IO rate from any initiator node **110** with which it interacts. Storage node **120** may choose a TCP window size based on a maximum QoS and worst case end-to-end latency, and thereby limit the demand from storage node **120**. Such a process may also limit the buffering needed to be performed by storage node **120**.

(134) In some examples, each of DPUs **117** may be capable of pipeline processing with zero stalls for data processing using tight integration between hardware and software. DPUs **117** may be capable of processing units of work (i.e., “work units”) in a run-to-completion model by a processor without any stalls or interruptions. A storage pipeline may be a set of work units, with each work unit executed on a different processor (or virtual processor) for processing the IOPs associated with a volume. An example pipeline of work units for a raw (i.e., ephemeral) volume may take the form of: TCP.Math.NVMe over Fabrics.Math.Raw Volume.Math.Backend NVMe Driver

(135) A durable volume, which provides additional functionality of log structuring and erasure coding/replication, may add additional stages to the pipeline. Parallel work units may also be executed in forks and joins. Similarly, additional features like encryption, compression, and deduplication might add more stages to the pipeline. Further details about work units and pipeline processing may be found in U.S. Pat. No. 10,841,245, filed Nov. 20, 2018, entitled “WORK UNIT STACK DATA STRUCTURES IN MULTIPLE CORE PROCESSOR SYSTEM FOR STREAM DATA PROCESSING,” and U.S. Pat. No. 10,929,175, entitled “SERVICE CHAINING HARDWARE ACCELERATORS WITHIN A DATA STREAM PROCESSING INTEGRATED CIRCUIT”, filed Nov. 21, 2018. These publications are hereby incorporated by reference.

(136) Generally, there is a state associated with most of the stages of a storage pipeline and so relying on locking procedures that require overhead and/or context switches is not optimal. Therefore, processor or virtual processor may use serialization (serializing the state updates on a given processor) to perform state updates. As a result, an allocation/placement algorithm may be used to place the storage pipelines of each volume among the processors. Such

allocation/placement may be performed by accounting for the QoS parameters of each volume. This allocation logic may take various factors into consideration in choosing processors (or DPUs **117**) for a storage pipeline (of a volume). The processors for the storage pipeline of a volume requiring a given QoS may be selected by converting the QoS into CPU load in processor cycles or IOPs and checking the current load on the processors.

(137) This allocation/placement problem is similar to the bin packing problem, and to solve this problem, the following considerations may be used:

(138) (1) Cluster Proximity: each of DPUs **117** may have multiple processing clusters with each cluster comprising a number of CPU cores and each core comprising of a number of VPs (Virtual Processors). Cluster proximity may be used when selecting VPs for a given pipeline to avoid any cache miss stalls and also get the benefit of latency in accessing the cluster's local resources.

(139) (2) Pick Mostly Loaded Cluster First: the utilization of a processing cluster may be maximized before allocating a different cluster.

(140) (3) Pick Least Loaded Core First: CPU cores may be selected within a cluster in a way that maximizes performance; this may require uniform spreading of the load.

(141) A cost table, which may be statically populated based on performance data, can be used to represent the VP cycles required to service a given stage of the pipeline for a single IO operation. The cost data for both read and write operations are available if the system supports independent configuration of read and write QoS, as described herein. Use the maximum of costs of read QoS and write QoS can be used for reserving VP cycles. The VP allocation algorithm for a given pipeline may first select the mostly loaded cluster where the pipeline will fit and then may select the least loaded core. Once a core is selected for a stage of the pipeline (where a given work unit handler in the pipeline runs), a VP within that core may be allocated based on the least loaded VP.

(142) Scheduling in the networking VPs (i.e., networking pool **202** of FIG. 2B) generally ensures no more work is admitted than what a given DPU **117** can handle. Further, VP allocation for storage pipelines (volumes) may ensure that pipelines are allocated without any conflict. Note that the pipeline VP allocation within a DPU may be initially performed at the volume creation time but may also be adjusted dynamically at the time a volume is active. The dynamic relocation of a volume's pipeline when it is becoming active may be performed pursuant to the following steps: 1. Start queueing incoming work at the TCP or NVMe over Fabrics layer. Eventually, this might back pressure the Initiator if relocation is not quick. 2. Make sure there is no pending work by flushing all stages of the volume pipeline. 3. Make sure all the completions are processed as well (wait for completions to match submissions). 4. Change the VP allocation of volume pipeline. 5. Start processing queued traffic, so the Initiator can resume if it was back-pressured.

(143) FIG. 6 is a flow diagram illustrating operations performed by an example storage cluster in accordance with one or more aspects of the present disclosure. FIG. 6 is described below within the context of storage cluster **102** of FIG. 1B. In other examples, operations described in FIG. 6 may be performed by one or more other components, modules, systems, or devices. Further, in other examples, operations described in connection with FIG. 6 may be merged, performed in a difference sequence, omitted, or may encompass additional operations not specifically illustrated or described.

(144) In the process illustrated in FIG. 6, and in accordance with one or more aspects of the present disclosure, storage cluster **102** may establish a quality of service standard for each of a plurality of tenants sharing a storage resource provided by the plurality of storage nodes in the storage cluster, wherein the quality of service standard includes a guaranteed allocation of the storage resource for each of the plurality of tenants and a maximum allocation of the storage resource for each of the plurality of tenants (**601**). For example, with reference to FIG. 1B, controller **130** may detect input that it determines corresponds to input about service level agreement that specifies, for one or more tenants of storage cluster **102**, a set of standards that outline performance, availability, or capacity expectations associated with services provided by storage cluster **102** to such tenants. In some

examples, the set of standards may specify a guaranteed allocation of performance, availability, capacity, or other metric or attribute for a given resource within storage cluster **102**. Further, the set of standards may specify a maximum allocation of performance, availability, capacity, or other metric or attribute for the given resource.

(145) Storage cluster **102** may allocate a volume of storage within the storage cluster by identifying a set of storage nodes to provide the storage resource for the volume of storage, wherein the set of storage nodes are a subset of the plurality of storage nodes (**602**). For example, again with reference to FIG. **1B**, storage services module **131** of controller **130** determines, based on input or otherwise, that a new volume (e.g., volume **121J**) is to be created for a specific tenant of storage cluster **102** (e.g., tenant “J”). Storage services module **131** further determines the type of volume to be allocated and attributes of volume **121J** (e.g., compressed, encrypted, durable). Storage services module **131** accesses data store **132** and uses a best fit algorithm to identify storage nodes **120** to be allocated to volume **121J**. In one specific example, storage services module **131** determines that new volume **121J** is to be allocated using resources from storage nodes **120A**, **120B**, and **120D**. Storage services module **131** causes controller **130** to allocate volume **121J**.

(146) Storage cluster **102** may schedule an operation to be performed by the set of storage nodes for the volume of storage (**603**). For example, still referring to FIG. **1B**, one or more of the storage nodes **120** allocated to volume **121J**, such as storage node **120A**, detects input that it determines corresponds to work associated with an operation to be performed on volume **121J**. Storage node **120A** further determines that the work requires use of one or more resources of DPU **117A** and/or storage node **120A** in order to complete the work (e.g., where storage node **120A** has been allocated to volume **121J**). Storage node **120A** enqueues the work and waits until sufficient credits are available for the resource(s) needed (NO path from **604**). Once sufficient resources are available, storage node **120A** dequeues the work (**605** and YES path from **604**). In connection with dequeuing the work, storage node **120A** enables DPU **117A** to execute and/or perform the work (**606**).

(147) Various examples, descriptions, attributes, features, and/or techniques in accordance with one or more aspects of this disclosure are outlined in the clauses below.

(148) Clause 1: A method of achieving guaranteed and adaptive quality of service (QoS) in a scale out and disaggregated storage cluster. Such a method may comprise one or more DPU-based storage target nodes with attached SSDs, zero or more DPU powered storage initiator nodes, and one or more storage controller nodes for storage management and orchestration (i.e., cluster services).

(149) Clause 2: The method of any other clause, wherein the storage orchestration system uses a best-fit algorithm for allocating the cluster resources (SSDs, DPUs, and network bandwidth) to the volumes based on QoS requirements and a current allocation of the existing volumes and other factors like fault zones and SSD wear leveling.

(150) Clause 3: The method of any other clause, wherein the incoming work is scheduled by using Excessive Weighted Round Robin (EWRR) algorithm that provides guaranteed QoS of each active volume and also allows each volume to get up to maximum QoS by sharing unused resources among the active volumes.

(151) Clause 4: The method of any other clause, where the internal processor resources of a DPU are managed as three separate groups for network, volume, and SSD driver functions in order to ensure that volume function becomes the first bottleneck.

(152) Clause 5: The method of any other clause, wherein multiple NVMe queues belonging to a volume are rate limited together using a single rate limiter.

(153) Clause 6: The method of any other clause, including allowing rate limiting in a hierarchical fashion to achieve QoS of an aggregated volume group at one level and QoS of individual volumes within a volume group at the next level.

(154) Clause 7: The method of any other clause, including segregating the new work from inter DPU to DPU traffic, including traffic received by a DPU for journal replication and inter DPU

traffic for durability, and prioritize the new work without being subject to any rate limiting or QoS policies in order to ensure that the minimum QoS is met even for volumes spanning multiple DPUs.

(155) Clause 8: The method of any other clause, including supporting and/or configuring independent QoS parameters for read-only and write-only workloads and providing the QoS for mixed workloads based on the configured read-only and write-only QoS parameters and the ratio of read to write traffic in the mixed workload.

(156) Clause 9: The method of any other clause, wherein the incoming work of volumes is rate limited to maximum QoS by restricting the number of NVMeoF Queues (NVMe/TCP Queues when using TCP as the transport) and queue depth of the queues based on the maximum QoS parameters.

(157) Clause 10: The method of any other clause, wherein the TCP Receive Window Size is configured based on the maximum QoS when using NVMe/TCP.

(158) Clause 11: The method of any other clause, wherein the internal resources of a DPU, like processing cycles (e.g., VP cycles), are allocated based on QoS requirements.

(159) Clause 12: The method of any other clause, where the storage software functionality in a DPU is realized as a pipeline of work unit handlers is dynamically moved across the processor resources when a new active volume contends with other active volumes for resources within a DPU.

(160) Clause 13: The method of any other clause, including allowing oversubscription of the QoS of the cluster and relocate the volumes across DPUs when a DPU cannot meet the QoS of all the active volumes.

(161) Clause 14: The method of any other clause, where the QoS degrades for all volumes adaptively when a cluster is oversubscribed.

(162) Clause 15: A system comprising means for performing any of the methods of any prior clause.

(163) Clause 16: A system or computing device comprising means for performing any of the methods of any prior clause.

(164) For processes, apparatuses, and other examples or illustrations described herein, including in any flowcharts or flow diagrams, certain operations, acts, steps, or events included in any of the techniques described herein can be performed in a different sequence, may be added, merged, or left out altogether (e.g., not all described acts or events are necessary for the practice of the techniques). Moreover, in certain examples, operations, acts, steps, or events may be performed concurrently, e.g., through multi-threaded processing, interrupt processing, or multiple processors, rather than sequentially. Further certain operations, acts, steps, or events may be performed automatically even if not specifically identified as being performed automatically. Also, certain operations, acts, steps, or events described as being performed automatically may be alternatively not performed automatically, but rather, such operations, acts, steps, or events may be, in some examples, performed in response to input or another event.

(165) The disclosures of all publications, patents, and patent applications referred to herein are each hereby incorporated by reference in their entireties. To the extent that any such disclosure material that is incorporated by reference conflicts with the instant disclosure, the instant disclosure shall control.

(166) For ease of illustration, only a limited number of devices (e.g., initiator nodes **110**, storage nodes **120**, controllers **130**, as well as others) are shown within the Figures and/or in other illustrations referenced herein. However, techniques in accordance with one or more aspects of the present disclosure may be performed with many more of such systems, components, devices, modules, and/or other items, and collective references to such systems, components, devices, modules, and/or other items may represent any number of such systems, components, devices, modules, and/or other items.

(167) The Figures included herein each illustrate at least one example implementation of an aspect of this disclosure. The scope of this disclosure is not, however, limited to such implementations. Accordingly, other example or alternative implementations of systems, methods or techniques described herein, beyond those illustrated in the Figures, may be appropriate in other instances. Such implementations may include a subset of the devices and/or components included in the Figures and/or may include additional devices and/or components not shown in the Figures.

(168) The detailed description set forth above is intended as a description of various configurations and is not intended to represent the only configurations in which the concepts described herein may be practiced. The detailed description includes specific details for the purpose of providing a sufficient understanding of the various concepts. However, these concepts may be practiced without these specific details. In some instances, well-known structures and components are shown in block diagram form in the referenced figures in order to avoid obscuring such concepts.

(169) Accordingly, although one or more implementations of various systems, devices, and/or components may be described with reference to specific Figures, such systems, devices, and/or components may be implemented in a number of different ways. For instance, one or more devices illustrated herein as separate devices may alternatively be implemented as a single device; one or more components illustrated as separate components may alternatively be implemented as a single component. Also, in some examples, one or more devices illustrated in the Figures herein as a single device may alternatively be implemented as multiple devices; one or more components illustrated as a single component may alternatively be implemented as multiple components. Each of such multiple devices and/or components may be directly coupled via wired or wireless communication and/or remotely coupled via one or more networks. Also, one or more devices or components that may be illustrated in various Figures herein may alternatively be implemented as part of another device or component not shown in such Figures. In this and other ways, some of the functions described herein may be performed via distributed processing by two or more devices or components.

(170) Further, certain operations, techniques, features, and/or functions may be described herein as being performed by specific components, devices, and/or modules. In other examples, such operations, techniques, features, and/or functions may be performed by different components, devices, or modules. Accordingly, some operations, techniques, features, and/or functions that may be described herein as being attributed to one or more components, devices, or modules may, in other examples, be attributed to other components, devices, and/or modules, even if not specifically described herein in such a manner.

(171) Although specific advantages have been identified in connection with descriptions of some examples, various other examples may include some, none, or all of the enumerated advantages. Other advantages, technical or otherwise, may become apparent to one of ordinary skill in the art from the present disclosure. Further, although specific examples have been disclosed herein, aspects of this disclosure may be implemented using any number of techniques, whether currently known or not, and accordingly, the present disclosure is not limited to the examples specifically described and/or illustrated in this disclosure.

(172) In one or more examples, the functions described may be implemented in hardware, software, firmware, or any combination thereof. If implemented in software, the functions may be stored, as one or more instructions or code, on and/or transmitted over a computer-readable medium and executed by a hardware-based processing unit. Computer-readable media may include computer-readable storage media, which corresponds to a tangible medium such as data storage media, or communication media including any medium that facilitates transfer of a computer program from one place to another (e.g., pursuant to a communication protocol). In this manner, computer-readable media generally may correspond to (1) tangible computer-readable storage media, which is non-transitory or (2) a communication medium such as a signal or carrier wave. Data storage media may be any available media that can be accessed by one or more computers or

one or more processors to retrieve instructions, code and/or data structures for implementation of the techniques described in this disclosure. A computer program product may include a computer-readable medium.

(173) By way of example, and not limitation, such computer-readable storage media can include RAM, ROM, EEPROM, or optical disk storage, magnetic disk storage, or other magnetic storage devices, flash memory, or any other medium that can be used to store desired program code in the form of instructions or data structures and that can be accessed by a computer. Also, any connection may properly be termed a computer-readable medium. For example, if instructions are transmitted from a website, server, or other remote source using a wired (e.g., coaxial cable, fiber optic cable, twisted pair) or wireless (e.g., infrared, radio, and microwave) connection, then the wired or wireless connection is included in the definition of medium. It should be understood, however, that computer-readable storage media and data storage media do not include connections, carrier waves, signals, or other transient media, but are instead directed to non-transient, tangible storage media.

(174) Instructions may be executed by one or more processors, such as one or more digital signal processors (DSPs), general purpose microprocessors, application specific integrated circuits (ASICs), field programmable logic arrays (FPGAs), or other equivalent integrated or discrete logic circuitry. Accordingly, the terms “processor” or “processing circuitry” as used herein may each refer to any of the foregoing structure or any other structure suitable for implementation of the techniques described. In addition, in some examples, the functionality described may be provided within dedicated hardware and/or software modules. Also, the techniques could be fully implemented in one or more circuits or logic elements.

(175) The techniques of this disclosure may be implemented in a wide variety of devices or apparatuses, including a wireless handset, a mobile or non-mobile computing device, a wearable or non-wearable computing device, an integrated circuit (IC) or a set of ICs (e.g., a chip set). Various components, modules, or units are described in this disclosure to emphasize functional aspects of devices configured to perform the disclosed techniques, but do not necessarily require realization by different hardware units. Rather, as described above, various units may be combined in a hardware unit or provided by a collection of interoperating hardware units, including one or more processors as described above, in conjunction with suitable software and/or firmware.

Claims

1. A method comprising: establishing, by a storage cluster having a plurality of storage nodes, a quality of service standard for each of a plurality of tenants sharing a storage resource, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, wherein the quality of service standards for the plurality of tenants allow for oversubscription where utilization of the storage resource can exceed 100%, and wherein the storage resource is provided by the plurality of storage nodes in the storage cluster; allocating, by the storage cluster and based on the quality of service standards, a plurality of volumes of storage within the storage cluster, wherein allocating a volume of the plurality of volumes of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes; upon detecting that utilization of the storage resource is above a predetermined threshold, moving one or more of the plurality of volumes to address a risk that the quality of service standards for the plurality of tenants are not satisfied; scheduling, by the storage cluster and based on the quality of service standards, operations to be performed by the set of storage nodes for a given volume of storage, wherein scheduling operations includes: enqueueing, within a first storage node in the set of storage nodes, a plurality of demands to use the storage resource provided by the first storage node; maintaining a credit count associated with the plurality of demands; periodically increasing the credit count over time; dequeuing, based on the increased

credit count, one of the plurality of demands to enable an operation to be performed by a data processing unit of the first storage node; and after the dequeuing, decreasing the credit count; and executing the enabled operation to be performed by the data processing unit of the first storage node associated with the dequeued one of the plurality of demands.

2. The method of claim 1, wherein establishing the quality of service standard includes: establishing, for a first operation type, a first guaranteed allocation of the storage resource for operations of the first operation type for the tenant; establishing, for the first operation type, a first maximum allocation of the storage resource for operations of the first operation type for the tenant; establishing, for a second operation type, a second guaranteed allocation of the storage resource for operations of the second operation type for the tenant; and establishing, for the second operation type, a second maximum allocation of the storage resource for operations of the second operation type for the tenant.

3. The method of claim 1, wherein establishing the quality of service standard includes: establishing, for read operations, a guaranteed allocation of the storage resource for read operations for the tenant and a maximum allocation of the storage resource for read operations for the tenant; and establishing, for write operations, a guaranteed allocation of the storage resource for write operations for the tenant and a maximum allocation of the storage resource for write operations for the tenant.

4. The method of claim 1, wherein the storage resource is a first storage resource, wherein each of the plurality of tenants share the first storage resource and a second storage resource in the storage cluster, and wherein establishing the quality of service standard includes: establishing a guaranteed allocation of the second storage resource for the tenant; and establishing a maximum allocation of the second storage resource for the tenant.

5. The method of claim 1, wherein the storage resource includes at least one of: storage capacity, input/output processing cycles, bandwidth, processor processing cycles, or accelerator resources.

6. The method of claim 1, wherein allocating the volume of the plurality of volumes of storage includes: maintaining a count associated with the storage resource across the plurality of storage nodes in the storage cluster; and adjusting the count based on identifying the set of storage nodes to provide the storage resource for the volume of storage.

7. The method of claim 1, the method further comprising: deallocating, by the storage cluster, the volume of the plurality of volumes of storage, wherein deallocating the volume includes making at least some resources associated with the set of storage nodes for the volume available to be allocated to a new volume of storage.

8. The method of claim 1, wherein the plurality of tenants includes a first tenant and a second tenant, the method further comprising: establishing a first guaranteed allocation of the storage resource for the first tenant, and a first maximum allocation of the storage resource for the first tenant; and establishing a second guaranteed allocation of the storage resource for the second tenant, and a second maximum allocation of the storage resource for the second tenant, wherein the first guaranteed allocation is different than the second guaranteed allocation.

9. The method of claim 8, wherein the first maximum allocation is different than the second maximum allocation.

10. A storage cluster comprising: a network; and a plurality of computing systems, at least one of the computing systems comprising one or more processors, each interconnected over the network, wherein the plurality of computing systems includes a plurality of storage nodes comprising one or more storage machines, and wherein the plurality of computing systems are collectively configured to: establish a quality of service standard for each of a plurality of tenants sharing a storage resource, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, wherein the quality of service standards for the plurality of tenants allow for oversubscription where utilization of the storage resource can exceed 100%, and wherein the storage resource is provided

by the plurality of storage nodes in the storage cluster; allocate, based on the quality of service standards, a plurality of volumes of storage within the storage cluster, wherein allocating a volume of the plurality of volumes of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes; upon detecting that utilization of the storage resource is above a predetermined threshold, move one or more of the plurality of volumes to address a risk that the quality of service standards for the plurality of tenants are not satisfied; schedule, based on the quality of service standards, operations to be performed by the set of storage nodes for a given volume of storage, wherein scheduling operations includes: enqueueing, within a first storage node in the set of storage nodes, a plurality of demands to use the storage resource provided by the first storage node; maintaining a credit count associated with the plurality of demands; periodically increasing the credit count over time; dequeuing, based on the increased credit count, one of the plurality of demands to enable an operation to be performed by a data processing unit of the first storage node; and after the dequeuing, decreasing the credit count; and execute the enabled operation to be performed by the data processing unit of the first storage node associated with the dequeued one of the plurality of demands.

11. The storage cluster of claim 10, wherein to establish the quality of service standard, the computing systems are further configured to: establish, for a first operation type, a first guaranteed allocation of the storage resource for operations of the first operation type for the tenant; establish, for the first operation type, a first maximum allocation of the storage resource for operations of the first operation type for the tenant; establish, for a second operation type, a second guaranteed allocation of the storage resource for operations of the second operation type for the tenant; and establish, for the second operation type, a second maximum allocation of the storage resource for operations of the second operation type for the tenant.

12. The storage cluster of claim 10, wherein to establish the quality of service standard, the computing systems are further configured to: establish, for read operations, a guaranteed allocation of the storage resource for read operations for the tenant and a maximum allocation of the storage resource for read operations for the tenant; and establish, for write operations, a guaranteed allocation of the storage resource for write operations for the tenant and a maximum allocation of the storage resource for write operations for the tenant.

13. The storage cluster of claim 10, wherein the resource is a first storage resource, wherein each of the plurality of tenants share the first storage resource and a second storage resource in the storage cluster, and wherein to establish the quality of service standard, the computing systems are further configured to: establish a guaranteed allocation of the second storage resource for the tenant; and establish a maximum allocation of the second storage resource for the tenant.

14. The storage cluster of claim 10, wherein the storage resource includes at least one of: storage capacity, input/output processing cycles, bandwidth, processor processing cycles, or accelerator resources.

15. The storage cluster of claim 10, wherein to allocate the volume of the plurality of volumes of storage, the computing systems are further configured to: maintain a count associated with the storage resource across the plurality of storage nodes in the storage cluster; and adjust the count based on identifying the set of storage nodes to provide the storage resource for the volume of storage.

16. The storage cluster of claim 10, wherein the computing systems are further configured to: deallocate the volume of the plurality of volumes of storage, wherein deallocating the volume includes making at least some resources associated with the set of storage nodes for the volume available to be allocated to a new volume.

17. The storage cluster of claim 10, wherein the plurality of tenants include a first tenant and a second tenant, and wherein to establish the quality of service standard, the computing systems are further configured to: establish a first guaranteed allocation of the storage resource for the first

tenant, and a first maximum allocation of the storage resource for the first tenant; and establish a second guaranteed allocation of the storage resource for the second tenant, and a second maximum allocation of the storage resource for the second tenant, wherein the first guaranteed allocation is different than the second guaranteed allocation.

18. A storage cluster comprising processing circuitry and a system for storing computing instructions, wherein the processing circuitry has access to the system for storing computing instructions and is configured to: establish a quality of service standard for each of a plurality of tenants sharing a storage resource, wherein the quality of service standard includes a guaranteed allocation of the storage resource for the tenant and a maximum allocation of the storage resource for the tenant, wherein the quality of service standards for the plurality of tenants allow for oversubscription where utilization of the storage resource can exceed 100%, and wherein the storage resource is provided by the plurality of storage nodes in the storage cluster; allocate, based on the quality of service standards, a plurality of volumes of storage within the storage cluster, wherein allocating a volume of the plurality of volumes of storage includes identifying a set of storage nodes to provide the storage resource for the volume of storage, and wherein the set of storage nodes are a subset of the plurality of storage nodes; upon detecting that utilization of the storage resource is above a predetermined threshold, move one or more of the plurality of volumes to address a risk that the quality of service standards for the plurality of tenants are not satisfied; schedule, based on the quality of service standard, operations to be performed by the set of storage nodes for a given volume of storage, wherein scheduling operations includes: enqueueing, within a first storage node in the set of storage nodes, a plurality of demands to use the storage resource provided by the first storage node; maintaining a credit count associated with the plurality of demands; periodically increasing the credit count over time; dequeuing, based on the increased credit count, one of the plurality of demands to enable an operation to be performed by a data processing unit of the first storage node; and after the dequeuing, decreasing the credit count; and execute the enabled operation to be performed by the data processing unit of the first storage node associated with the dequeued one of the plurality of demands.
