

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411755
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Mankala; Santosh et al.

Testing integration applications for systems

Abstract

An integration application including a plurality of components may be provided. The components may provide connections between systems. The integration application may be configured to allow the systems to exchange data. It may be determined that a new version of a first one of the components is a candidate for release. Prior to release of the new version of the first component, the integration application may be tested using the new version of the first component and an existing version of the first component. An instance of the integration application may be automatically caused to be upgraded to include the new version of the first component.

Inventors: Mankala; Santosh (Mountain House, CA), Yanzon; Carlos Santiago (Buenos Aires, AR), Sabino; Jose (Buenos Aires, AR), McGrath; Keith (Chicago, IL), Goyal; Nimit (Naperville, IL), Surana; Prince (Milpitas, CA), Melakhessou; Wassim (Suisun City, CA), Glaser; Scott (Dublin, CA), Balireddy; Siddharth (San Francisco, CA), Kale; Tarun (San Francisco, CA)

Applicant: Salesforce, Inc. (San Francisco, CA)

Family ID: 90359224

Assignee: Salesforce, Inc. (San Francisco, CA)

Appl. No.: 17/935822

Filed: September 27, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240104003 A1	Mar. 28, 2024

Publication Classification

Int. Cl.: G06F11/3668 (20250101)

U.S. Cl.:

CPC **G06F11/368** (20130101); **G06F11/3684** (20130101); **G06F11/3688** (20130101);

Field of Classification Search

CPC: G06F (11/368); G06F (11/3684); G06F (11/3688); G06F (11/3692)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8739126	12/2013	Glaser et al.	N/A	N/A
8868981	12/2013	Glaser	714/45	G06F 11/3684
8903943	12/2013	Glaser et al.	N/A	N/A
9563680	12/2016	Mankala et al.	N/A	N/A
10114737	12/2017	Ramakrishna et al.	N/A	N/A
10261890	12/2018	Gandham et al.	N/A	N/A
10503495	12/2018	Kolhe	N/A	G06F 8/36
10846210	12/2019	Singh	N/A	G06F 8/71
2014/0005974	12/2013	Vyas et al.	N/A	N/A
2014/0026122	12/2013	Markande	717/124	G06F 11/3698
2020/0218641	12/2019	Schwenker	N/A	G16H 40/40

OTHER PUBLICATIONS

Ruian Duan; Automating Patching of Vulnerable Open-Source Software Versions in Application Binaries; NDSS; 15 pages; retrieved on May 23, 2025 (Year: 2019). cited by examiner

Primary Examiner: Soug; S.

Assistant Examiner: Luu; Cuong V

Attorney, Agent or Firm: Kwan & Olynick LLP

Background/Summary

FIELD OF TECHNOLOGY

(1) This patent document relates generally to on-demand software applications, and more specifically to application programming interfaces associated with such on-demand software applications.

BACKGROUND

(2) “Cloud computing” services provide shared resources, applications, and information to computers and other devices upon request. In cloud computing environments, services can be provided by one or more servers accessible over the Internet rather than installing software locally on in-house computer systems. Users can interact with cloud computing services to undertake a wide range of tasks.

(3) Cloud-based services may include software applications that are implemented in a distributed manner, and have multiple services and microservices distributed across different components of a

cloud computing environment. Such services and microservices may be implemented to perform particular tasks and functions of the software applications. Application programming interfaces (APIs) may be used to facilitate communication between such services and applications. Conventional applications remain limited in their ability to instantiate and update APIs in such large-scale and heterogenous computing environments.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The included drawings are for illustrative purposes and serve only to provide examples of possible structures and operations for the disclosed inventive systems, apparatus, methods and computer program products for content synchronization. These drawings in no way limit any changes in form and detail that may be made by one skilled in the art without departing from the spirit and scope of the disclosed implementations.
- (2) FIG. 1 illustrates an example of an arrangement of components in a distributed system, configured in accordance with one or more implementations.
- (3) FIG. 2 illustrates an example of a method for testing integration applications for systems in accordance with one or more implementations
- (4) FIG. 3 illustrates example architecture of an integration application, according to one or more implementations.
- (5) FIG. 4 illustrates an example of an environment in which integration applications may be provided, according to one or more implementations.
- (6) FIG. 5 illustrates an example of a series of tests being performed on an integration application, according to one or more implementations.
- (7) FIG. 6 illustrates another example of a method for testing integration applications, according to one or more implementations.
- (8) FIG. 7 shows a block diagram of an example of an environment that includes an on-demand database service configured in accordance with some implementations.
- (9) FIG. 8A shows a system diagram of an example of architectural components of an on-demand database service environment, configured in accordance with some implementations.
- (10) FIG. 8B shows a system diagram further illustrating an example of architectural components of an on-demand database service environment, in accordance with some implementations.
- (11) FIG. 9 illustrates one example of a computing device.

DETAILED DESCRIPTION

(12) Software systems may be implemented in a distributed context in which they are provided as an on-demand service that is accessible to multiple users. Such distributed software systems may be implemented using various services that are dispersed throughout the distributed computing environment. Such services may use function calls to issue requests and responses to each other as well as with components of other applications and systems. Interfaces, such as APIs, may be used to facilitate such communication and provide connections between heterogeneous systems. An “integration application,” made up of components that provide such connections, may be implemented to allow such complex systems to exchange data with other complex systems. In complex distributed software systems with multiple instantiations of various services, the number of such components in any given integration application may become large. By way of example, Briljula Forensics makes software that communicates with hundreds of different crime and forensic systems and databases around the world. Briljula Forensics uses an integration platform to create an integration application to allow communication between these databases and their own software. Since many of these crime and forensic systems use different data formats and communication protocols, the integration application has many different components.

(13) Conventional integration applications may be prone to failure as new versions of components are continually released. Traditionally, when a new version of a particular component is released, the integration application instance may be automatically upgraded to the use new version of the component. As discussed above, integration applications may include a number of interconnected components that interact with one another. Accordingly, if there is an issue with the new version of the component, the entire integration application may start failing. Returning to the example of the previous paragraph, unbeknownst to Erik an authorized administrator in charge of maintaining Briljula Forensics' software, the integration platform updates a version of a connector component that the Briljula Forensics integration application uses to provide an interface between their internal software and external software used by a variety of law enforcement organizations. Unfortunately, the connector component interacts with several other components, which each, in turn, interact with other components. The updated version of the connector component does not interact properly with other existing components, creating an error which propagates via other component interactions. As a result, the Briljula Forensics integration application fails.

(14) By contrast, the disclosed techniques may be implemented to provide robust integration applications that are resistant to failure as new versions of components are released. For example, a customer of an integration platform may create configurable testing integrated with a release cycle. Similarly, the customer can tweak components of the integration application using the testing infrastructure disclosed herein without having to entirely re-build the integration application, saving significant time. By way of illustration, returning to the above example, when the Briljula Forensics integration application is uploaded, Erik may select that tests be run automatically before any upgrade. Tests may then be run automatically on the integration application using both old and new versions of each component of the integration as soon as a release candidate is available. As discussed below, if the integration application passes the tests using the new version of the component, the integration application may be automatically updated to include the new version of the component. Alternatively, if there is a failure, a report may be generated, and the customer may automatically be added to an exclusion list.

(15) In some implementations, the disclosed techniques may be deployed to create a variety of technical improvements in addition to those described above. For example, as new components are added to an integration application and new versions are released, there may be an exponentially increasing number of different combinations of components and versions of these components. By storing the results of tests of integration applications, as discussed in further detail below, an integration platform may store millions of test cases and unique scenarios. Storing data related to such an exponentially increasing amount of scenarios may provide a deep understanding of complex structures involved in integration applications with many components.

(16) In yet another example, the disclosed techniques allow for full automation of onboarding integration applications, with the system reporting errors as components of these integration applications are changed. By way of illustration, with an automated testing structure in place, there is no need for any human oversight when a new integration application is created by a customer of an integration platform. Similarly, as components of the integration application are changed, errors may be reported automatically as discussed herein.

(17) FIG. 1 illustrates an example of an arrangement of components in a distributed system, configured in accordance with one or more implementations. As will be discussed in greater detail below, a system, such as system **100**, may be implemented to facilitate testing of integration applications, as disclosed herein. More specifically, one or more components of system **100** may be configured to make a software integration application available as an on-demand service, and additional components are configured to facilitate automation of testing of integration applications that may be used for such an on-demand service.

(18) Accordingly, system **100** includes one or more client machines, which may also be referred to herein as client devices, such as client machine **102**. In various implementations, client machine

102 is a computing device accessible by a user. For example, client machine **102** may be a desktop computer, a laptop computer, a mobile computing device such as a smartphone, or any other suitable computing device. Accordingly, client machine **102** includes one or more input and display devices, and is communicatively coupled to communications network **130**, such as the internet. In various implementations, client machine **102** comprises one or more processors configured to execute one or more applications that may utilize a user interface. Accordingly, a user may request and view various different display screens associated with such applications via client machine **102**. In various implementations, a user interface may be used to present the display screen to the user, as well as receive one or more inputs from the user. In some implementations, the user interface may utilize a web browser executed on client machine **102** or may be a standalone locally executed application. Moreover, such user interfaces may be used to access on-demand services and software applications, as will be discussed in greater detail below.

(19) In various implementations, system **100** further includes one or more servers configured to provide a computing platform, such as computing platform **112**, and various client devices may be communicatively coupled to computing platform **112**. In various implementations, computing platform **112** is configured to include software and hardware that provides an environment for the execution of an integration application. As will be discussed in greater detail below, computing platform **112** may implement an integration platform and may be configured to provide and test integration applications with a variety of components each with a variety of versions.

(20) Computing platform **112** may include one or more processors and memory configured execute a software application. Accordingly, computing platform **112** may be configured to store program code and settings for a particular application, and may also be configured to execute the code. Moreover, computing platform **112** may be in communication with numerous client devices, and may implement the integration application in a distributed manner. In some implementations, computing platform **112** is further configured to generate and serve webpages that may be viewed by a user via one or more devices, such as client machine **102**. Accordingly, computing platform **112** is configured to provide a web-based interface between a user of client machine **102** and an integration application that is deployed in a distributed environment. In some implementations, computing platform **112** is coupled to database system **114** which may be configured to store various application data and data associated with webpages served by computing platform **112**, and thus may provide local storage for computing platform **112**.

(21) FIG. 2 illustrates an example of a method for testing integration applications for systems in accordance with one or more implementations. FIG. 2 is discussed in the context of FIGS. 3-6. FIG. 3 illustrates example architecture of an integration application, according to one or more implementations. FIG. 4 illustrates an example of an environment in which integration applications may be provided, according to one or more implementations. FIG. 5 illustrates an example of a series of tests being performed on an integration application, according to one or more implementations. FIG. 6 illustrates another example of a method for testing integration applications, according to one or more implementations.

(22) At **202** of FIG. 2, an integration application (e.g., integration application **300** of FIG. 3) may be provided. The integration application **300** may integrate systems and allow these systems to exchange data, as described above. Such systems may encompass a range of constituents such as applications, services, any combination of hardware or software components, etc.

(23) The integration application **300** may have M versions **302**. As discussed above, the integration application **302** may be made up of components **304(1)-(N)**. The components **304(1)-(N)** may be any type of application constituent. For example, such components may perform a variety of functions such as: asynchronous processing of parts of a flow or subflow, batch processing of messages, initiating subflows, logging, setting payloads, transforming messages, etc.

(24) In some implementations, as discussed above, the components **304(1)-(N)** may provide a variety of connections between systems. By way of example, such components may include

connectors that facilitate the integration of integration applications with any external resources, such as third-party APIs or standard integration protocols. Such connectors may enable connectivity to SaaS and on-premises applications, systems, and services.

(25) Also or alternatively, the components **304(1)-(N)** may include modules that act internally in the integration application without providing a connection to a particular endpoint. For instance, such modules may perform a variety of functions such as aggregating values, compressing data, using Java features, processing and extracting data from XML documents, etc.

(26) In some implementations, each component **304(1)-(N)** may have any arbitrary number **M.sub.i** of versions **308(1)-(M.sub.i)**. Together, all of the **N** components **308(1)-(N)** make up the integration application **300**.

(27) In some implementations, the integration application **300** may be implemented along with a testing framework that allows authorized users to create, configure, and maintain automated tests. Such a testing framework may allow users to create and build tests by writing in a variety of languages such as Java code or a coding language native to the integration platform. Such a testing framework may be implemented to perform a variety of functions such as verifying message processor calls, mock message processing, creating mock outbound endpoints, disabling outbound endpoints, disabling flow inbound connectors, disabling inbound endpoints, debugging tests created using the framework, viewing coverage reports, etc. By way of example, an authorized user affiliated with an organization (e.g., Erik of Briljula Forensics, as described in the above example) may create any arbitrary number **Q** tests **312(1)-(Q)** to test the integration application **300** in a variety of contexts.

(28) Referring now to FIG. 4, an integration application (e.g., integration application **300** of FIG. 3) may be provided in environment **400** of FIG. 4 via exchange **404**. The exchange **404** may be a service where integration applications are managed by an integration platform (e.g., the computing platform **112** of FIG. 1).

(29) In some implementations, an integration application (e.g., the integration application **300** of FIG. 3) may run on a container **408** of FIG. 4. The container **408** may be provided by the integration platform via nexus **412**. The nexus **412** may be managed by the integration platform and may be a location from which the container **408** is obtained. Since the integration platform provides the container **408**, the container may be where testing of integration applications occurs, as described further below. Additionally, the integration platform may update the containers **408** and automatically verify that integration applications are running on the latest available container.

(30) Also or alternatively, the containers **408** may store dependencies such that when integration applications are tested using different permutations of components (as described further below); result service **416** causes results of these tests to be stored in result database **420**. As described above, having a large repository of such test results may confer benefits beyond those related to each individual integration application being tested. For example, certain combinations of components of integration applications may fail under a variety of circumstances. Therefore, maintaining the results database **420** may allow the integration platform to have data pertaining to millions of different integration application scenarios, providing a broad understanding of combinations of components.

(31) Returning to FIG. 2, at **204**, it may be determined that a new version of a component of the integration application is a candidate for release. For instance, returning to the above example, the integration platform (e.g., the computing platform **112** of FIG. 1) may automatically determine that the new version of the connector component used by the Briljula Forensics integration application is a candidate for release.

(32) At **206** of FIG. 2, the integration application may be tested. By way of illustration, returning to the above example, the integration platform (e.g., the computing platform **112** of FIG. 1) may automatically begin testing the Briljula Forensics integration application in response to determining that the new version of the connector component is a candidate for release.

(33) Such testing may occur in a variety of manners. By way of example, a test may be automatically run on integration an application using a first permutation of versions of each component of the integration application. After automatically testing the integration application using the first permutation of versions of each component of the integration application, the integration application may be changed to use a different permutation of versions of each component and the integration application may be tested using the different permutation of versions. This process may be repeated until every possible permutation of versions of the integration application has been tested.

(34) By way of illustration, referring now to the general example depicted in FIG. 3, the integration application **300** of FIG. 3 has N components **304(1)-(N)** each having M.sub.i versions **308**. Test **1 312(1)** may be run on a first permutation of versions of components **304(1)-(N)**. For example, the test **1 312(1)** may be run on the integration application **300** using version 1 of each component **304**. The test **1 312(1)** may then be run on the integration application **300** using version 2 of component **1 304(1)** and version 1 of all other components **304**. This process may be repeated until each and every permutation of versions **308** of components **304** have been tested using the test **1 312(1)**.

(35) Since there are Q tests, this process may be repeated where each permutation of versions **308** of components **304** may be tested for each of test **1 312(1)** through test Q **312(Q)**.

(36) Referring now to a simplified example depicted in FIG. 5, integration application **500** has two components **504(1)** and **504(2)**. The component **504(1)** has versions **508(1) 1-X** (X being any arbitrary number). The component **504(2)** has versions **508(2) 1-Y** (Y being any arbitrary number). For example, test **1 512(1)** may first be run on the integration application **500** using version 1 of component **1 504(1)** and component **2 504(2)**. The test **1 512(1)** may then be run on the integration application **500** using version 2 of component **1 504(1)** and version 1 of component **2 504(2)**. The test **1 512(1)** may then continue to be run on the integration application **500** using a different version **508(1)** of component **1 504(1)** and version 1 of component **2 504(2)** until all X versions **508(1)** of the component **1 504(1)** have been tested with version 1 of component **2 504(2)**.

(37) Once all X versions **508(1)** of the component **1 504(1)** have been tested with version 1 of component **2 504(2)**, the test **1 512(1)** may then be run on the integration application **500** X further times using each of the X versions **508(1)** of the component **1 504(1)** along with version 2 of component **2 504(2)**. This process may be repeated until the test **1 512(1)** has been run on the integration application **500** using all X versions **508(1)** of the component **1 504(1)** along with all Y versions **508(2)** of component **2 504(2)**.

(38) The above process may be repeated for test **2 512(2)** such that test **2 512(2)** has been run on the integration application **500** using all X versions **508(1)** of the component **1 504(1)** along with all Y versions **508(2)** of component **2 504(2)**. Similarly, the above process may be repeated for test **3 512(3)** such that test **3 512(3)** has been run on the integration application **500** using all X versions **508(1)** of the component **1 504(1)** along with all Y versions **508(2)** of component **2 504(2)**.

(39) Returning to FIG. 2, at **208**, the integration application may be caused to be upgraded. By way of example, if the test(s) performed at **206** determine that the integration application **300** of FIG. 3 executes properly with the new version of the component, then the integration application may be automatically upgraded to include the new version of the component.

(40) Also or alternatively, if the integration application fails any of the test(s) at **206**, a failure report may be automatically generated responsive to the failure. By way of illustration, responsive to a failure associated with a test of the integration application **300** of FIG. 3, the integration platform (e.g., the computing platform **112** of FIG. 1) may automatically generate and store a report of the results of the test in the results database **420** of FIG. 4. In the event of this failure, the integration application **300** of FIG. 3 may be automatically added to an exclusion list.

(41) Referring now to FIG. 6, method **600** may be performed on integration application **602**, which may be any type of integration application described herein such as the integration application **300** of FIG. 3. The integration application **602** of FIG. 6 may have one or more components that have

been changed and the method **600** may be performed on integration application **602** to test integration application **602**, as described above.

(42) At **604**, it may be determined whether the integration application **602** is self-hosted or hosted by a computing platform (e.g., the computing platform **112** of FIG. 1). If the integration application **602** is self-hosted, method **600** may proceed to **608** where a container may be obtained (e.g., the container **408** of FIG. 4, as described above). In this case, if the integration application **602** is self-hosted, running on a customer infrastructure separate from the computing platform associated with the testing infrastructure, the integration application **602** may be automatically on-boarded to run on the container, at which point the method **600** may proceed to **620**, and the container may be run.

(43) Alternatively, if it is determined at **604** of FIG. 6 that the integration application **602** is hosted by the computing platform (e.g., the computing platform **112** of FIG. 1), the method **600** may proceed to **612** at which point the integration application **602** may be uploaded to an exchange (e.g., the exchange **404** of FIG. 4). In this case, if the integration application **602** is hosted by the computing platform, requirements may be configured by the computing platform, and the customer may not need to provide any configuration information. After requirements are configured, the method **600** may proceed to **620**, and the container may be run.

(44) At **624**, it may be determined if a version of a component of the integration application **602** is defined. If it is determined that the version of the component of the integration application **602** is not defined, the method **600** may proceed to **628** and the latest version of the component of the integration application **602** may be used in testing at **632**, as discussed below. Alternatively, if it is determined that the version of the component of the integration application **602** is defined, the method **600** may proceed directly to **632** at which point testing may be executed using the defined version of the integration application **602**.

(45) At **632** of FIG. 6, the integration application **602** is tested using the techniques described above in the context of **206** of FIG. 2. Once the integration application **602** is tested, the method **600** proceeds to **636**.

(46) If it is determined at **636** that there were not any problems associated with the results of the tests, the method **600** proceeds to **640** at which point the payload may be sent to a service API endpoint such that the integration application may be automatically updated to include the tested versions of the components. Additionally, the test results may be sent to a database (e.g., result service **416** of FIG. 4 may cause the results to be stored in the result database **420**, as described above.)

(47) Alternatively, if it is determined at **636** that there were problems with the tests, the method **600** proceeds to **640** at which point logs may be created and sent to a database (e.g., result service **416** of FIG. 4 may cause the results to be stored in the result database **420**.)

(48) FIG. 7 shows a block diagram of an example of an environment **710** that includes an on-demand database service configured in accordance with some implementations. Environment **710** may include user systems **712**, network **714**, database system **716**, processor system **717**, application platform **718**, network interface **720**, tenant data storage **722**, tenant data **723**, system data storage **724**, system data **725**, program code **726**, process space **728**, User Interface (UI) **730**, Application Program Interface (API) **732**, PL/SOQL **734**, save routines **736**, application setup mechanism **738**, application servers **750-1** through **750-N**, system process space **752**, tenant process spaces **754**, tenant management process space **760**, tenant storage space **762**, user storage **764**, and application metadata **766**. Some of such devices may be implemented using hardware or a combination of hardware and software and may be implemented on the same physical device or on different devices. Thus, terms such as “data processing apparatus,” “machine,” “server” and “device” as used herein are not limited to a single hardware device, but rather include any hardware and software configured to provide the described functionality.

(49) An on-demand database service, implemented using system **716**, may be managed by a database service provider. Some services may store information from one or more tenants into

tables of a common database image to form a multi-tenant database system (MTS). As used herein, each MTS could include one or more logically and/or physically connected servers distributed locally or across one or more geographic locations. Databases described herein may be implemented as single databases, distributed databases, collections of distributed databases, or any other suitable database system. A database image may include one or more database objects. A relational database management system (RDBMS) or a similar system may execute storage and retrieval of information against these objects.

(50) In some implementations, the application platform **718** may be a framework that allows the creation, management, and execution of applications in system **716**. Such applications may be developed by the database service provider or by users or third-party application developers accessing the service. Application platform **718** includes an application setup mechanism **738** that supports application developers' creation and management of applications, which may be saved as metadata into tenant data storage **722** by save routines **736** for execution by subscribers as one or more tenant process spaces **754** managed by tenant management process **760** for example. Invocations to such applications may be coded using PL/SOQL **734** that provides a programming language style interface extension to API **732**. A detailed description of some PL/SOQL language implementations is discussed in commonly assigned U.S. Pat. No. 7,730,478, titled METHOD AND SYSTEM FOR ALLOWING ACCESS TO DEVELOPED APPLICATIONS VIA A MULTI-TENANT ON-DEMAND DATABASE SERVICE, by Craig Weissman, issued on Jun. 1, 2010, and hereby incorporated by reference in its entirety and for all purposes. Invocations to applications may be detected by one or more system processes. Such system processes may manage retrieval of application metadata **766** for a subscriber making such an invocation. Such system processes may also manage execution of application metadata **766** as an application in a virtual machine.

(51) In some implementations, each application server **750** may handle requests for any user associated with any organization. A load balancing function (e.g., an F5 Big-IP load balancer) may distribute requests to the application servers **750** based on an algorithm such as least-connections, round robin, observed response time, etc. Each application server **750** may be configured to communicate with tenant data storage **722** and the tenant data **723** therein, and system data storage **724** and the system data **725** therein to serve requests of user systems **712**. The tenant data **723** may be divided into individual tenant storage spaces **762**, which can be either a physical arrangement and/or a logical arrangement of data. Within each tenant storage space **762**, user storage **764** and application metadata **766** may be similarly allocated for each user. For example, a copy of a user's most recently used (MRU) items might be stored to user storage **764**. Similarly, a copy of MRU items for an entire tenant organization may be stored to tenant storage space **762**. A UI **730** provides a user interface and an API **732** provides an application programming interface to system **716** resident processes to users and/or developers at user systems **712**.

(52) System **716** may implement a web-based integration system. For example, in some implementations, system **716** may include application servers configured to implement and execute software applications. The application servers may be configured to provide related data, code, forms, web pages and other information to and from user systems **712**. Additionally, the application servers may be configured to store information to, and retrieve information from a database system. Such information may include related data, objects, and/or Webpage content. With a multi-tenant system, data for multiple tenants may be stored in the same physical database object in tenant data storage **722**, however, tenant data may be arranged in the storage medium(s) of tenant data storage **722** so that data of one tenant is kept logically separate from that of other tenants. In such a scheme, one tenant may not access another tenant's data, unless such data is expressly shared.

(53) Several elements in the system shown in FIG. 7 include conventional, well-known elements that are explained only briefly here. For example, user system **712** may include processor system **712A**, memory system **712B**, input system **712C**, and output system **712D**. A user system **712** may be implemented as any computing device(s) or other data processing apparatus such as a mobile

phone, laptop computer, tablet, desktop computer, or network of computing devices. User system **712** may run an internet browser allowing a user (e.g., a subscriber of an MTS) of user system **712** to access, process and view information, pages and applications available from system **716** over network **714**. Network **714** may be any network or combination of networks of devices that communicate with one another, such as any one or any combination of a LAN (local area network), WAN (wide area network), wireless network, or other appropriate configuration.

(54) The users of user systems **712** may differ in their respective capacities, and the capacity of a particular user system **712** to access information may be determined at least in part by “permissions” of the particular user system **712**. As discussed herein, permissions generally govern access to computing resources such as data objects, components, and other entities of a computing system, such as an integration system, a social networking system, and/or a CRM database system. “Permission sets” generally refer to groups of permissions that may be assigned to users of such a computing environment. For instance, the assignments of users and permission sets may be stored in one or more databases of System **716**. Thus, users may receive permission to access certain resources. A permission server in an on-demand database service environment can store criteria data regarding the types of users and permission sets to assign to each other. For example, a computing device can provide to the server data indicating an attribute of a user (e.g., geographic location, industry, role, level of experience, etc.) and particular permissions to be assigned to the users fitting the attributes. Permission sets meeting the criteria may be selected and assigned to the users. Moreover, permissions may appear in multiple permission sets. In this way, the users can gain access to the components of a system.

(55) In some an on-demand database service environments, an Application Programming Interface (API) may be configured to expose a collection of permissions and their assignments to users through appropriate network-based services and architectures, for instance, using Simple Object Access Protocol (SOAP) Web Service and Representational State Transfer (REST) APIs.

(56) In some implementations, a permission set may be presented to an administrator as a container of permissions. However, each permission in such a permission set may reside in a separate API object exposed in a shared API that has a child-parent relationship with the same permission set object. This allows a given permission set to scale to millions of permissions for a user while allowing a developer to take advantage of joins across the API objects to query, insert, update, and delete any permission across the millions of possible choices. This makes the API highly scalable, reliable, and efficient for developers to use.

(57) In some implementations, a permission set API constructed using the techniques disclosed herein can provide scalable, reliable, and efficient mechanisms for a developer to create tools that manage a user's permissions across various sets of access controls and across types of users. Administrators who use this tooling can effectively reduce their time managing a user's rights, integrate with external systems, and report on rights for auditing and troubleshooting purposes. By way of example, different users may have different capabilities with regard to accessing and modifying application and database information, depending on a user's security or permission level, also called authorization. In systems with a hierarchical role model, users at one permission level may have access to applications, data, and database information accessible by a lower permission level user, but may not have access to certain applications, database information, and data accessible by a user at a higher permission level.

(58) As discussed above, system **716** may provide on-demand database service to user systems **712** using an MTS arrangement. By way of example, one tenant organization may be a company that employs a sales force where each salesperson uses system **716** to manage their sales process. Thus, a user in such an organization may maintain contact data, leads data, customer follow-up data, performance data, goals and progress data, etc., all applicable to that user's personal sales process (e.g., in tenant data storage **722**). In this arrangement, a user may manage his or her sales efforts and cycles from a variety of devices, since relevant data and applications to interact with (e.g.,

access, view, modify, report, transmit, calculate, etc.) such data may be maintained and accessed by any user system **712** having network access.

(59) When implemented in an MTS arrangement, system **716** may separate and share data between users and at the organization-level in a variety of manners. For example, for certain types of data each user's data might be separate from other users' data regardless of the organization employing such users. Other data may be organization-wide data, which is shared or accessible by several users or potentially all users form a given tenant organization. Thus, some data structures managed by system **716** may be allocated at the tenant level while other data structures might be managed at the user level. Because an MTS might support multiple tenants including possible competitors, the MTS may have security protocols that keep data, applications, and application use separate. In addition to user-specific data and tenant-specific data, system **716** may also maintain system-level data usable by multiple tenants or other data. Such system-level data may include industry reports, news, postings, and the like that are sharable between tenant organizations.

(60) In some implementations, user systems **712** may be client systems communicating with application servers **750** to request and update system-level and tenant-level data from system **716**. By way of example, user systems **712** may send one or more queries requesting data of a database maintained in tenant data storage **722** and/or system data storage **724**. An application server **750** of system **716** may automatically generate one or more SQL statements (e.g., one or more SQL queries) that are designed to access the requested data. System data storage **724** may generate query plans to access the requested data from the database.

(61) The database systems described herein may be used for a variety of database applications. By way of example, each database can generally be viewed as a collection of objects, such as a set of logical tables, containing data fitted into predefined categories. A “table” is one representation of a data object, and may be used herein to simplify the conceptual description of objects and custom objects according to some implementations. It should be understood that “table” and “object” may be used interchangeably herein. Each table generally contains one or more data categories logically arranged as columns or fields in a viewable schema. Each row or record of a table contains an instance of data for each category defined by the fields. For example, a CRM database may include a table that describes a customer with fields for basic contact information such as name, address, phone number, fax number, etc. Another table might describe a purchase order, including fields for information such as customer, product, sale price, date, etc. In some multi-tenant database systems, standard entity tables might be provided for use by all tenants. For CRM database applications, such standard entities might include tables for case, account, contact, lead, and opportunity data objects, each containing pre-defined fields. It should be understood that the word “entity” may also be used interchangeably herein with “object” and “table”.

(62) In some implementations, tenants may be allowed to create and store custom objects, or they may be allowed to customize standard entities or objects, for example by creating custom fields for standard objects, including custom index fields. Commonly assigned U.S. Pat. No. 7,779,039, titled CUSTOM ENTITIES AND FIELDS IN A MULTI-TENANT DATABASE SYSTEM, by Weissman et al., issued on Aug. 17, 2010, and hereby incorporated by reference in its entirety and for all purposes, teaches systems and methods for creating custom objects as well as customizing standard objects in an MTS. In certain implementations, for example, all custom entity data rows may be stored in a single multi-tenant physical table, which may contain multiple logical tables per organization. It may be transparent to customers that their multiple “tables” are in fact stored in one large table or that their data may be stored in the same table as the data of other customers.

(63) FIG. **8A** shows a system diagram of an example of architectural components of an on-demand database service environment **800**, configured in accordance with some implementations. A client machine located in the cloud **804** may communicate with the on-demand database service environment via one or more edge routers **808** and **812**. A client machine may include any of the examples of user systems **712** described above. The edge routers **808** and **812** may communicate

with one or more core switches **820** and **824** via firewall **816**. The core switches may communicate with a load balancer **828**, which may distribute server load over different pods, such as the pods **840** and **844** by communication via pod switches **832** and **836**. The pods **840** and **844**, which may each include one or more servers and/or other computing resources, may perform data processing and other operations used to provide on-demand services. Components of the environment may communicate with a database storage **856** via a database firewall **848** and a database switch **852**.

(64) Accessing an on-demand database service environment may involve communications transmitted among a variety of different components. The environment **800** is a simplified representation of an actual on-demand database service environment. For example, some implementations of an on-demand database service environment may include anywhere from one to many devices of each type. Additionally, an on-demand database service environment need not include each device shown, or may include additional devices not shown, in FIGS. **8A** and **8B**.

(65) The cloud **804** refers to any suitable data network or combination of data networks, which may include the Internet. Client machines located in the cloud **804** may communicate with the on-demand database service environment **800** to access services provided by the on-demand database service environment **800**. By way of example, client machines may access the on-demand database service environment **800** to retrieve, store, edit, and/or process information.

(66) In some implementations, the edge routers **808** and **812** route packets between the cloud **804** and other components of the on-demand database service environment **800**. The edge routers **808** and **812** may employ the Border Gateway Protocol (BGP). The edge routers **808** and **812** may maintain a table of IP networks or 'prefixes', which designate network reachability among autonomous systems on the internet.

(67) In one or more implementations, the firewall **816** may protect the inner components of the environment **800** from internet traffic. The firewall **816** may block, permit, or deny access to the inner components of the on-demand database service environment **800** based upon a set of rules and/or other criteria. The firewall **816** may act as one or more of a packet filter, an application gateway, a stateful filter, a proxy server, or any other type of firewall.

(68) In some implementations, the core switches **820** and **824** may be high-capacity switches that transfer packets within the environment **800**. The core switches **820** and **824** may be configured as network bridges that quickly route data between different components within the on-demand database service environment. The use of two or more core switches **820** and **824** may provide redundancy and/or reduced latency.

(69) In some implementations, communication between the pods **840** and **844** may be conducted via the pod switches **832** and **836**. The pod switches **832** and **836** may facilitate communication between the pods **840** and **844** and client machines, for example via core switches **820** and **824**. Also or alternatively, the pod switches **832** and **836** may facilitate communication between the pods **840** and **844** and the database storage **856**. The load balancer **828** may distribute workload between the pods, which may assist in improving the use of resources, increasing throughput, reducing response times, and/or reducing overhead. The load balancer **828** may include multilayer switches to analyze and forward traffic.

(70) In some implementations, access to the database storage **856** may be guarded by a database firewall **848**, which may act as a computer application firewall operating at the database application layer of a protocol stack. The database firewall **848** may protect the database storage **856** from application attacks such as structure query language (SQL) injection, database rootkits, and unauthorized information disclosure. The database firewall **848** may include a host using one or more forms of reverse proxy services to proxy traffic before passing it to a gateway router and/or may inspect the contents of database traffic and block certain content or database requests. The database firewall **848** may work on the SQL application level atop the TCP/IP stack, managing applications' connection to the database or SQL management interfaces as well as intercepting and enforcing packets traveling to or from a database network or application interface.

(71) In some implementations, the database storage **856** may be an on-demand database system shared by many different organizations. The on-demand database service may employ a single-tenant approach, a multi-tenant approach, a virtualized approach, or any other type of database approach. Communication with the database storage **856** may be conducted via the database switch **852**. The database storage **856** may include various software components for handling database queries. Accordingly, the database switch **852** may direct database queries transmitted by other components of the environment (e.g., the pods **840** and **844**) to the correct components within the database storage **856**.

(72) FIG. **8B** shows a system diagram further illustrating an example of architectural components of an on-demand database service environment, in accordance with some implementations. The pod **844** may be used to render services to user(s) of the on-demand database service environment **800**. The pod **844** may include one or more content batch servers **864**, content search servers **868**, query servers **882**, file servers **886**, access control system (ACS) servers **880**, batch servers **884**, and app servers **888**. Also, the pod **844** may include database instances **890**, quick file systems (QFS) **892**, and indexers **894**. Some or all communication between the servers in the pod **844** may be transmitted via the switch **836**.

(73) In some implementations, the app servers **888** may include a framework dedicated to the execution of procedures (e.g., programs, routines, scripts) for supporting the construction of applications provided by the on-demand database service environment **800** via the pod **844**. One or more instances of the app server **888** may be configured to execute all or a portion of the operations of the services described herein.

(74) In some implementations, as discussed above, the pod **844** may include one or more database instances **890**. A database instance **890** may be configured as an MTS in which different organizations share access to the same database, using the techniques described above. Database information may be transmitted to the indexer **894**, which may provide an index of information available in the database **890** to file servers **886**. The QFS **892** or other suitable filesystem may serve as a rapid-access file system for storing and accessing information available within the pod **844**. The QFS **892** may support volume management capabilities, allowing many disks to be grouped together into a file system. The QFS **892** may communicate with the database instances **890**, content search servers **868** and/or indexers **894** to identify, retrieve, move, and/or update data stored in the network file systems (NFS) **896** and/or other storage systems.

(75) In some implementations, one or more query servers **882** may communicate with the NFS **896** to retrieve and/or update information stored outside of the pod **844**. The NFS **896** may allow servers located in the pod **844** to access information over a network in a manner similar to how local storage is accessed. Queries from the query servers **882** may be transmitted to the NFS **896** via the load balancer **828**, which may distribute resource requests over various resources available in the on-demand database service environment **800**. The NFS **896** may also communicate with the QFS **892** to update the information stored on the NFS **896** and/or to provide information to the QFS **892** for use by servers located within the pod **844**.

(76) In some implementations, the content batch servers **864** may handle requests internal to the pod **844**. These requests may be long-running and/or not tied to a particular customer, such as requests related to log mining, cleanup work, and maintenance tasks. The content search servers **868** may provide query and indexer functions such as functions allowing users to search through content stored in the on-demand database service environment **800**. The file servers **886** may manage requests for information stored in the file storage **898**, which may store information such as documents, images, basic large objects (BLOBS), etc. The query servers **882** may be used to retrieve information from one or more file systems. For example, the query system **882** may receive requests for information from the app servers **888** and then transmit information queries to the NFS **896** located outside the pod **844**. The ACS servers **880** may control access to data, hardware resources, or software resources called upon to render services provided by the pod **844**.

The batch servers **884** may process batch jobs, which are used to run tasks at specified times. Thus, the batch servers **884** may transmit instructions to other servers, such as the app servers **888**, to trigger the batch jobs.

(77) While some of the disclosed implementations may be described with reference to a system having an application server providing a front end for an on-demand database service capable of supporting multiple tenants, the disclosed implementations are not limited to multi-tenant databases nor deployment on application servers. Some implementations may be practiced using various database architectures such as ORACLE®, DB2® by IBM and the like without departing from the scope of present disclosure.

(78) FIG. **9** illustrates one example of a computing device. According to various embodiments, a system **900** suitable for implementing embodiments described herein includes a processor **901**, a memory module **903**, a storage device **905**, an interface **911**, and a bus **915** (e.g., a PCI bus or other interconnection fabric.) System **900** may operate as variety of devices such as an application server, a database server, or any other device or service described herein. Although a particular configuration is described, a variety of alternative configurations are possible. The processor **901** may perform operations such as those described herein. Instructions for performing such operations may be embodied in the memory **903**, on one or more non-transitory computer readable media, or on some other storage device. Various specially configured devices can also be used in place of or in addition to the processor **901**. The interface **911** may be configured to send and receive data packets over a network. Examples of supported interfaces include, but are not limited to: Ethernet, fast Ethernet, Gigabit Ethernet, frame relay, cable, digital subscriber line (DSL), token ring, Asynchronous Transfer Mode (ATM), High-Speed Serial Interface (HSSI), and Fiber Distributed Data Interface (FDDI). These interfaces may include ports appropriate for communication with the appropriate media. They may also include an independent processor and/or volatile RAM. A computer system or computing device may include or communicate with a monitor, printer, or other suitable display for providing any of the results mentioned herein to a user.

(79) Any of the disclosed implementations may be embodied in various types of hardware, software, firmware, computer readable media, and combinations thereof. For example, some techniques disclosed herein may be implemented, at least in part, by computer-readable media that include program instructions, state information, etc., for configuring a computing system to perform various services and operations described herein. Examples of program instructions include both machine code, such as produced by a compiler, and higher-level code that may be executed via an interpreter. Instructions may be embodied in any suitable language such as, for example, Apex, Java, Python, C++, C, HTML, any other markup language, JavaScript, ActiveX, VBScript, or Perl. Examples of computer-readable media include, but are not limited to: magnetic media such as hard disks and magnetic tape; optical media such as flash memory, compact disk (CD) or digital versatile disk (DVD); magneto-optical media; and other hardware devices such as read-only memory (“ROM”) devices and random-access memory (“RAM”) devices. A computer-readable medium may be any combination of such storage devices.

(80) In the foregoing specification, various techniques and mechanisms may have been described in singular form for clarity. However, it should be noted that some embodiments include multiple iterations of a technique or multiple instantiations of a mechanism unless otherwise noted. For example, a system uses a processor in a variety of contexts but can use multiple processors while remaining within the scope of the present disclosure unless otherwise noted. Similarly, various techniques and mechanisms may have been described as including a connection between two entities. However, a connection does not necessarily mean a direct, unimpeded connection, as a variety of other entities (e.g., bridges, controllers, gateways, etc.) may reside between the two entities.

(81) In the foregoing specification, reference was made in detail to specific embodiments including one or more of the best modes contemplated by the inventors. While various implementations have

been described herein, it should be understood that they have been presented by way of example only, and not limitation. For example, some techniques and mechanisms are described herein in the context of on-demand computing environments that include MTSs. However, the techniques of disclosed herein apply to a wide variety of computing environments. Particular embodiments may be implemented without some or all of the specific details described herein. In other instances, well known process operations have not been described in detail in order to avoid unnecessarily obscuring the disclosed techniques. Accordingly, the breadth and scope of the present application should not be limited by any of the implementations described herein, but should be defined only in accordance with the claims and their equivalents.

Claims

1. A method comprising: providing, via a computing platform, an integration application comprising a plurality of components, the components providing connections between systems, the integration application configured to allow the systems to exchange data; determining whether the integration application is self-hosted or hosted by the computing platform; uploading the integration application to an exchange according to a result of the determining whether the integration application is self-hosted or hosted by the computing platform; determining, using one or more processors, that a new version of a first component of the components is a candidate for release; prior to the release of the new version of the first component, and responsive to determining that the new version of the first component is the candidate for the release, testing, using the one or more processors, the integration application using the new version of the first component; automatically causing, based on results of the testing the integration application, an instance of the integration application to be upgraded to include the new version of the first component; running a further test on the integration application; and responsive to a failure associated with the further test: automatically generating and storing a report of results of the further test; and automatically adding the integration application to an exclusion list.
2. The method of claim 1, wherein the testing the integration application comprises: automatically testing the integration application using each permutation of versions of each component of the integration application.
3. The method of claim 2, wherein the automatically testing the integration application using each permutation of versions of each component of the integration application comprises: automatically testing the integration application using a first permutation of versions of each component of the integration application; and after the automatically testing the integration application using the first permutation of versions of each component of the integration application, automatically changing the integration application to use a different permutation of versions of each component of the integration application and testing the integration application using the different permutation of versions of each component of the integration application until all permutations of versions have been tested.
4. The method of claim 1, wherein the computing platform is provided to a plurality of organizations via an on-demand computing environment, the integration application being implemented by a first one of the organizations.
5. The method of claim 4, wherein the testing of the integration application is one of a plurality of tests configurable by authorized users associated with an organization implementing the computing platform.
6. The method of claim 5, further comprising: maintaining a results database, the results database storing the results for the testing of the integration application and storing further results of further tests of further integration applications associated with the plurality of organizations.
7. The method of claim 6, wherein the integration application is configured to run on a container associated with the computing platform regardless of whether the first organization uses a first

infrastructure associated with the computing platform or the first organization uses a second infrastructure native to the first organization.

8. A computing platform comprising one or more processors and a memory, the computing platform being configured to cause: providing, via a computing platform, an integration application comprising a plurality of components, the components providing connections between systems, the integration application configured to allow the systems to exchange data; determining whether the integration application is self-hosted or hosted by the computing platform; uploading the integration application to an exchange according to a result of the determining whether the integration application is self-hosted or hosted by the computing platform; determining, using the one or more processors, that a new version of a first component of the components is a candidate for release; prior to the release of the new version of the first component, and responsive to determining that the new version of the first component is the candidate for the release, testing, using the one or more processors, the integration application using the new version of the first component and an existing version of the first component; automatically causing, based on results of the testing the integration application, an instance of the integration application to be upgraded to include the new version of the first component; running a further test on the integration application; and responsive to a failure associated with the further test: automatically generating and storing a report of results of the further test; and automatically adding the integration application to an exclusion list.

9. The computing platform of claim 8, wherein the testing the integration application comprises: automatically testing the integration application using each permutation of versions of each component of the integration application.

10. The computing platform of claim 9, wherein the automatically testing the integration application using each permutation of versions of each component of the integration application comprises: automatically testing the integration application using a first permutation of versions of each component of the integration application; and after the automatically testing the integration application using the first permutation of versions of each component of the integration application, automatically changing the integration application to use a different permutation of versions of each component of the integration application and testing the integration application using the different permutation of versions of each component of the integration application until all permutations of versions have been tested.

11. The computing platform of claim 8, wherein the computing platform is provided to a plurality of organizations via an on-demand computing environment, the integration application being implemented by a first one of the organizations.

12. The computing platform of claim 11, wherein the testing of the integration application is one of a plurality of tests configurable by authorized users associated with an organization implementing the computing platform.

13. The computing platform of claim 12, the computing platform further configured to cause: maintaining a results database, the results database storing the results for the testing of the integration application and storing further results of further tests of further integration applications associated with the plurality of organizations.

14. The computing platform of claim 13, wherein the integration application is configured to run on a container associated with the computing platform regardless of whether the first organization uses a first infrastructure associated with the computing platform or the first organization uses a second infrastructure native to the first organization.

15. A non-transitory computer-readable medium comprising computer-readable program code executable by one or more processors, the program code comprising instructions configured to cause the one or more processors to perform a method comprising: providing, via a computing platform, an integration application comprising a plurality of components, the components providing connections between systems, the integration application configured to allow the systems

to exchange data; determining whether the integration application is self-hosted or hosted by the computing platform; uploading the integration application to an exchange according to a result of the determining whether the integration application is self-hosted or hosted by the computing platform; determining, using the one or more processors, that a new version of a first component of the components is a candidate for release; prior to the release of the new version of the first component, and responsive to determining that the new version of the first component is the candidate for release, testing, using the one or more processors, the integration application using the new version of the first component and an existing version of the first component; automatically causing, based on results of the testing the integration application, an instance of the integration application to be upgraded to include the new version of the first component; running a further test on the integration application; and responsive to a failure associated with the further test: automatically generating and storing a report of results of the further test; and automatically adding the integration application to an exclusion list.

16. The non-transitory computer-readable medium of claim 15, wherein the testing the integration application comprises: automatically testing the integration application using each permutation of versions of each component of the integration application.

17. The non-transitory computer-readable medium of claim 16, wherein the automatically testing the integration application using each permutation of versions of each component of the integration application comprises: automatically testing the integration application using a first permutation of versions of each component of the integration application; and after the automatically testing the integration application using the first permutation of versions of each component of the integration application, automatically changing the integration application to use a different permutation of versions of each component of the integration application and testing the integration application using the different permutation of versions of each component of the integration application until all permutations of versions have been tested.
