

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411766
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Zbiciak; Joseph

Streaming engine with variable stream template format

Abstract

A streaming engine employed in a digital data processor specifies a fixed read only data stream defined by plural nested loops. An address generator produces address of data elements for the nested loops. A steam head register stores data elements next to be supplied to functional units for use as operands. A stream template specifies loop count and loop dimension for each nested loop. A format definition field in the stream template specifies the number of loops and the stream template bits devoted to the loop counts and loop dimensions. This permits the same bits of the stream template to be interpreted differently enabling trade off between the number of loops supported and the size of the loop counts and loop dimensions.

Inventors:	Zbiciak; Joseph (San Jose, CA)
Applicant:	TEXAS INSTRUMENTS INCORPORATED (Dallas, TX)
Family ID:	62561705
Assignee:	Texas Instruments Incorporated (Dallas, TX)
Appl. No.:	18/594091
Filed:	March 04, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240202123 A1	Jun. 20, 2024

Related U.S. Application Data

continuation parent-doc US 17972675 20221025 US 11921636 child-doc US 18594091
continuation parent-doc US 17146576 20210112 US 11481327 20221025 child-doc US 17972675
continuation parent-doc US 16459210 20190701 US 10891231 20210112 child-doc US 17146576
continuation parent-doc US 15384487 20161220 US 10339057 20190702 child-doc US 16459210

Publication Classification

Int. Cl.: **G06F12/08** (20160101); **G06F9/30** (20180101); **G06F9/345** (20180101); **G06F9/38** (20180101); **G06F12/0815** (20160101); **G06F12/0862** (20160101); **G06F12/0875** (20160101); **G06F12/0897** (20160101); G06F9/32 (20180101); G06F15/80 (20060101)

U.S. Cl.:

CPC **G06F12/0815** (20130101); **G06F9/3001** (20130101); **G06F9/30036** (20130101); **G06F9/30038** (20230801); **G06F9/30047** (20130101); **G06F9/30072** (20130101); **G06F9/3012** (20130101); **G06F9/3013** (20130101); **G06F9/30145** (20130101); **G06F9/345** (20130101); **G06F9/3822** (20130101); **G06F9/383** (20130101); **G06F9/3853** (20130101); **G06F9/3887** (20130101); **G06F12/0862** (20130101); **G06F12/0875** (20130101); **G06F12/0897** (20130101); G06F9/30065 (20130101); G06F9/325 (20130101); G06F15/8007 (20130101); G06F2212/452 (20130101); G06F2212/454 (20130101); G06F2212/6026 (20130101); G06F2212/621 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5640578	12/1996	Balmer et al.	N/A	N/A
6671799	12/2002	Parthasarathy	N/A	N/A
6829696	12/2003	Balmer et al.	N/A	N/A
6832296	12/2003	Hooker	N/A	N/A
7926046	12/2010	Halambi	717/149	G06F 8/4452
8108652	12/2011	Hui	N/A	N/A
9170811	12/2014	Mudawar	N/A	G06F 9/30149
9606803	12/2016	Anderson et al.	N/A	N/A
9740484	12/2016	Gopal et al.	N/A	N/A
10109370	12/2017	Bacha	N/A	G11C 29/38
2004/0261067	12/2003	Komatsu et al.	N/A	N/A
2005/0015537	12/2004	Asaad et al.	N/A	N/A
2005/0071615	12/2004	Goodrich	N/A	N/A
2005/0273769	12/2004	Eichenberger et al.	N/A	N/A
2005/0283589	12/2004	Matsuo	N/A	N/A
2006/0212654	12/2005	Balakrishnan	N/A	N/A
2006/0212685	12/2005	Raghavan	712/E9.046	G06F 9/3824
2008/0141013	12/2007	Klima et al.	N/A	N/A
2009/0113191	12/2008	Hall	712/241	G06F 9/325
2009/0138661	12/2008	Lauterbach	712/E9.055	G06F 9/3017
2009/0189896	12/2008	Jiao et al.	N/A	N/A
2011/0029962	12/2010	Nuzman et al.	N/A	N/A
2012/0185714	12/2011	Chung	712/42	G06F 9/381
2013/0185540	12/2012	Hung et al.	N/A	N/A

2013/0339700	12/2012	Blasco-Allue	712/241	G06F 9/325
2014/0188961	12/2013	Plotnikov et al.	N/A	N/A
2014/0281369	12/2013	Jha	N/A	N/A
2015/0019840	12/2014	Anderson et al.	N/A	N/A
2015/0227374	12/2014	Blasco	712/240	G06F 9/3844
2017/0308381	12/2016	Zbiciak	N/A	N/A
2017/0344375	12/2016	Zhu et al.	N/A	N/A
2018/0173532	12/2017	Zbiciak	N/A	N/A

Primary Examiner: Coleman; Eric

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation of U.S. patent application Ser. No. 17/972,675, filed Oct. 25, 2022, which is a continuation of U.S. patent application Ser. No. 17/146,576, filed Jan. 12, 2021, now U.S. Pat. No. 11,481,327, which is a continuation of U.S. patent application Ser. No. 16/459,210, filed Jul. 1, 2019, now U.S. Pat. No. 10,891,231, which is a continuation of U.S. patent application Ser. No. 15/384,487, filed Dec. 20, 2016, now U.S. Pat. No. 10,339,057, each of which is incorporated herein by reference in its entirety. (2) This application is an improvement over U.S. patent application Ser. No. 14/331,986, filed Jul. 15, 2014, now U.S. Pat. No. 9,606,803, which claims priority to U.S. Provisional Patent Application No. 61/846,148, filed Jul. 15, 2013, each of which is incorporated herein by reference in its entirety.

TECHNICAL FIELD

(1) The technical field of this invention is digital data processing and more specifically control of streaming engine used for operand fetching.

BACKGROUND

(2) Modern digital signal processors (DSP) faces multiple challenges. Workloads continue to increase, requiring increasing bandwidth. Systems on a chip (SOC) continue to grow in size and complexity. Memory system latency severely impacts certain classes of algorithms. As transistors get smaller, memories and registers become less reliable. As software stacks get larger, the number of potential interactions and errors becomes larger.

(3) Memory bandwidth and scheduling are a problem for digital signal processors operating on real-time data. Digital signal processors operating on real-time data typically receive an input data stream, perform a filter function on the data stream (such as encoding or decoding) and output a transformed data stream. The system is called real-time because the application fails if the transformed data stream is not available for output when scheduled. Typical video encoding requires a predictable but non-sequential input data pattern. Often the corresponding memory accesses are difficult to achieve within available address generation and memory access resources. A typical application requires memory access to load data registers in a data register file and then supply to functional units which perform the data processing.

SUMMARY

(4) This invention is a streaming engine employed in a digital signal processor. A fixed data stream sequence is specified by storing corresponding parameters in a control register. The data stream includes plural nested loops. Once started the data stream is read only and cannot be written. A functional unit using the stream data has a first instruction type that only reads the data and a second instruction type both reads the data and causes the streaming engine to advance the stream. This generally corresponds to the needs of a real-time filtering operation.

- (5) The streaming engine includes an address generator which produces address of data elements and a stream head register which stores data elements next to be supplied to functional units for use as operands. Each of the plural nested loops has a specified loop count and loop dimension. A stream template specifies loop count and loop dimension for each nested loop. A format definition field in the stream template specifies the number of loops and the stream template bits devoted to the loop counts and loop dimensions. This permits the same bits of the stream template to be interpreted differently enabling tradeoff between the number of loops supported and the size of the loop counts and loop dimensions
- (6) The preferred embodiment includes two independently defined data streams. The two data streams may be independently read or read/advances by a set of very long instruction word (VLIW) functional units.
-

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) These and other aspects of this invention are illustrated in the drawings, in which:
- (2) FIG. 1 illustrates a dual scalar/vector datapath processor according to one embodiment of this invention;
- (3) FIG. 2 illustrates the registers and functional units in the dual scalar/vector datapath processor illustrated in FIG. 1;
- (4) FIG. 3 illustrates a global scalar register file;
- (5) FIG. 4 illustrates a local scalar register file shared by arithmetic functional units;
- (6) FIG. 5 illustrates a local scalar register file shared by multiply functional units;
- (7) FIG. 6 illustrates local scalar register file of shared by load/store units;
- (8) FIG. 7 illustrates a global vector register file;
- (9) FIG. 8 illustrates a predicate register file;
- (10) FIG. 9 illustrates a local vector register file shared by arithmetic functional units;
- (11) FIG. 10 illustrates a local vector register file shared by multiply and correlation functional units;
- (12) FIG. 11 illustrates pipeline phases of a central processing unit according to a preferred embodiment of this invention;
- (13) FIG. 12 illustrates sixteen instructions of a single fetch packet;
- (14) FIG. 13 illustrates an example of the instruction coding of instructions used by this invention;
- (15) FIG. 14 illustrates the bit coding of a condition code extension slot 0;
- (16) FIG. 15 illustrates the bit coding of a condition code extension slot 1;
- (17) FIG. 16 illustrates the bit coding of a constant extension slot 0;
- (18) FIG. 17 is a partial block diagram illustrating constant extension;
- (19) FIG. 18 illustrates the carry control for SIMD operations according to this invention;
- (20) FIG. 19 illustrates a conceptual view of the streaming engines of this invention;
- (21) FIG. 20 illustrates a first example of lane allocation in a vector;
- (22) FIG. 21 illustrates a second example of lane allocation in a vector;
- (23) FIG. 22 illustrates a basic two dimensional stream;
- (24) FIG. 23 illustrates the order of elements within the example stream of FIG. 21;
- (25) FIG. 24 illustrates extracting a smaller rectangle from a larger rectangle;
- (26) FIG. 25 illustrates how the streaming engine would fetch the stream of this example with a transposition granularity of 4 bytes;
- (27) FIG. 26 illustrates how the streaming engine would fetch the stream of this example with a transposition granularity of 8 bytes;
- (28) FIG. 27 illustrates the details of streaming engine of this invention;

- (29) FIG. **28** illustrates a stream template register of this invention;
- (30) FIG. **29** illustrates sub-field definitions of the flags field of the stream template register of this invention;
- (31) FIG. **30** illustrates a first alternate stream template register of this invention as specified by a dimension format field;
- (32) FIG. **31** illustrates a second alternate stream template register of this invention as specified by a dimension format field;
- (33) FIG. **32** illustrates a third alternate stream template register of this invention as specified by a dimension format field;
- (34) FIG. **33** illustrates a fourth alternate stream template register of this invention as specified by a dimension format field;
- (35) FIG. **34** illustrates a fifth alternate stream template register of this invention as specified by a dimension format field;
- (36) FIG. **35** illustrates a loop count selection circuit for selecting stream definition template bits for loop count for the plural nested loops as controlled by the dimensional format field;
- (37) FIG. **36** illustrates a loop dimension selection circuit for selecting stream definition template bits for loop dimension for the plural nested loops as controlled by the dimensional format field;
- (38) FIG. **37** illustrates an example of adder control word circuit which generates an adder control word for loop address generators controlling the selection of linear address mode and circular address mode;
- (39) FIG. **38** is a partial schematic diagram illustrating address generation for the nested loops of this invention;
- (40) FIG. **39** illustrates a partial schematic diagram showing the streaming engine supply of operands of this invention.

DETAILED DESCRIPTION

- (41) FIG. **1** illustrates a dual scalar/vector datapath processor according to a preferred embodiment of this invention. Processor **100** includes separate level one instruction cache (L1I) **121** and level one data cache (L1D) **123**. Processor **100** includes a level two combined instruction/data cache (L2) **130** that holds both instructions and data. FIG. **1** illustrates connection between level one instruction cache **121** and level two combined instruction/data cache **130** (bus **142**). FIG. **1** illustrates connection between level one data cache **123** and level two combined instruction/data cache **130** (bus **145**). In the preferred embodiment of processor **100** level two combined instruction/data cache **130** stores both instructions to back up level one instruction cache **121** and data to back up level one data cache **123**. In the preferred embodiment level two combined instruction/data cache **130** is further connected to higher level cache and/or main memory in a manner not illustrated in FIG. **1**. In the preferred embodiment central processing unit core **110**, level one instruction cache **121**, level one data cache **123** and level two combined instruction/data cache **130** are formed on a single integrated circuit. This signal integrated circuit optionally includes other circuits.
- (42) Central processing unit core **110** fetches instructions from level one instruction cache **121** as controlled by instruction fetch unit **111**. Instruction fetch unit **111** determines the next instructions to be executed and recalls a fetch packet sized set of such instructions. The nature and size of fetch packets are further detailed below. As known in the art, instructions are directly fetched from level one instruction cache **121** upon a cache hit (if these instructions are stored in level one instruction cache **121**). Upon a cache miss (the specified instruction fetch packet is not stored in level one instruction cache **121**), these instructions are sought in level two combined cache **130**. In the preferred embodiment the size of a cache line in level one instruction cache **121** equals the size of a fetch packet. The memory locations of these instructions are either a hit in level two combined cache **130** or a miss. A hit is serviced from level two combined cache **130**. A miss is serviced from a higher level of cache (not illustrated) or from main memory (not illustrated). As is known in the

art, the requested instruction may be simultaneously supplied to both level one instruction cache **121** and central processing unit core **110** to speed use.

(43) In the preferred embodiment of this invention, central processing unit core **110** includes plural functional units to perform instruction specified data processing tasks. Instruction dispatch unit **112** determines the target functional unit of each fetched instruction. In the preferred embodiment central processing unit **110** operates as a very long instruction word (VLIW) processor capable of operating on plural instructions in corresponding functional units simultaneously. Preferably a compiler organizes instructions in execute packets that are executed together. Instruction dispatch unit **112** directs each instruction to its target functional unit. The functional unit assigned to an instruction is completely specified by the instruction produced by a compiler. The hardware of central processing unit core **110** has no part in this functional unit assignment. In the preferred embodiment instruction dispatch unit **112** may operate on plural instructions in parallel. The number of such parallel instructions is set by the size of the execute packet. This will be further detailed below.

(44) One part of the dispatch task of instruction dispatch unit **112** is determining whether the instruction is to execute on a functional unit in scalar datapath side A **115** or vector datapath side B **116**. An instruction bit within each instruction called the s bit determines which datapath the instruction controls. This will be further detailed below.

(45) Instruction decode unit **113** decodes each instruction in a current execute packet. Decoding includes identification of the functional unit performing the instruction, identification of registers used to supply data for the corresponding data processing operation from among possible register files and identification of the register destination of the results of the corresponding data processing operation. As further explained below, instructions may include a constant field in place of one register number operand field. The result of this decoding is signals for control of the target functional unit to perform the data processing operation specified by the corresponding instruction on the specified data.

(46) Central processing unit core **110** includes control registers **114**. Control registers **114** store information for control of the functional units in scalar datapath side A **115** and vector datapath side B **116** in a manner not relevant to this invention. This information could be mode information or the like.

(47) The decoded instructions from instruction decode **113** and information stored in control registers **114** are supplied to scalar datapath side A **115** and vector datapath side B **116**. As a result functional units within scalar datapath side A **115** and vector datapath side B **116** perform instruction specified data processing operations upon instruction specified data and store the results in an instruction specified data register or registers. Each of scalar datapath side A **115** and vector datapath side B **116** includes plural functional units that preferably operate in parallel. These will be further detailed below in conjunction with FIG. 2. There is a datapath **117** between scalar datapath side A **115** and vector datapath side B **116** permitting data exchange.

(48) Central processing unit core **110** includes further non-instruction based modules. Emulation unit **118** permits determination of the machine state of central processing unit core **110** in response to instructions. This capability will typically be employed for algorithmic development. Interrupts/exceptions unit **119** enable central processing unit core **110** to be responsive to external, asynchronous events (interrupts) and to respond to attempts to perform improper operations (exceptions).

(49) Central processing unit core **110** includes streaming engine **125**. Streaming engine **125** supplies two data streams from predetermined addresses typically cached in level two combined cache **130** to register files of vector datapath side B. This provides controlled data movement from memory (as cached in level two combined cache **130**) directly to functional unit operand inputs. This is further detailed below.

(50) FIG. 1 illustrates exemplary data widths of busses between various parts. Level one instruction

cache **121** supplies instructions to instruction fetch unit **111** via bus **141**. Bus **141** is preferably a 512-bit bus. Bus **141** is unidirectional from level one instruction cache **121** to central processing unit core **110**. Level two combined cache **130** supplies instructions to level one instruction cache **121** via bus **142**. Bus **142** is preferably a 512-bit bus. Bus **142** is unidirectional from level two combined cache **130** to level one instruction cache **121**.

(51) Level one data cache **123** exchanges data with register files in scalar datapath side A **115** via bus **143**. Bus **143** is preferably a 64-bit bus. Level one data cache **123** exchanges data with register files in vector datapath side B **116** via bus **144**. Bus **144** is preferably a 512-bit bus. Busses **143** and **144** are illustrated as bidirectional supporting both central processing unit core **110** data reads and data writes. Level one data cache **123** exchanges data with level two combined cache **130** via bus **145**. Bus **145** is preferably a 512-bit bus. Bus **145** is illustrated as bidirectional supporting cache service for both central processing unit core **110** data reads and data writes.

(52) Level two combined cache **130** supplies data of a first data stream to streaming engine **125** via bus **146**. Bus **146** is preferably a 512-bit bus. Streaming engine **125** supplies data of this first data stream to functional units of vector datapath side B **116** via bus **147**. Bus **147** is preferably a 512-bit bus. Level two combined cache **130** supplies data of a second data stream to streaming engine **125** via bus **148**. Bus **148** is preferably a 512-bit bus. Streaming engine **125** supplies data of this second data stream to functional units of vector datapath side B **116** via bus **149**. Bus **149** is preferably a 512-bit bus. Busses **146**, **147**, **148** and **149** are illustrated as unidirectional from level two combined cache **130** to streaming engine **125** and to vector datapath side B **116** in accordance with the preferred embodiment of this invention.

(53) In the preferred embodiment of this invention, both level one data cache **123** and level two combined cache **130** may be configured as selected amounts of cache or directly addressable memory in accordance with U.S. Pat. No. 6,606,686 entitled UNIFIED MEMORY SYSTEM ARCHITECTURE INCLUDING CACHE AND DIRECTLY ADDRESSABLE STATIC RANDOM ACCESS MEMORY.

(54) FIG. 2 illustrates further details of functional units and register files within scalar datapath side A **115** and vector datapath side B **116**. Scalar datapath side A **115** includes global scalar register file **211**, L1/S1 local register file **212**, M1/N1 local register file **213** and D1/D2 local register file **214**. Scalar datapath side A **115** includes L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226**. Vector datapath side B **116** includes global vector register file **231**, L2/S2 local register file **232**, M2/N2/C local register file **233** and predicate register file **234**. Vector datapath side B **116** includes L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, C unit **245**, and P unit **246**. There are limitations upon which functional units may read from or write to which register files. These will be detailed below.

(55) Scalar datapath side A **115** includes L1 unit **221**. L1 unit **221** generally accepts two 64-bit operands and produces one 64-bit result. The two operands are each recalled from an instruction specified register in either global scalar register file **211** or L1/S1 local register file **212**. L1 unit **221** preferably performs the following instruction selected operations: 64-bit add/subtract operations; 32-bit min/max operations; 8-bit Single Instruction Multiple Data (SIMD) instructions such as sum of absolute value, minimum and maximum determinations; circular min/max operations; and various move operations between register files. The result may be written into an instruction specified register of global scalar register file **211**, L1/S1 local register file **212**, M1/N1 local register file **213**, or D1/D2 local register file **214**.

(56) Scalar datapath side A **115** includes S1 unit **222**. S1 unit **222** generally accepts two 64-bit operands and produces one 64-bit result. The two operands are each recalled from an instruction specified register in either global scalar register file **211** or L1/S1 local register file **212**. S1 unit **222** preferably performs the same type operations as L1 unit **221**. There optionally may be slight variations between the data processing operations supported by L1 unit **221** and S1 unit **222**. The result may be written into an instruction specified register of global scalar register file **211**, L1/S1

local register file **212**, M1/N1 local register file **213**, or D1/D2 local register file **214**.

(57) Scalar datapath side A **115** includes M1 unit **223**. M1 unit **223** generally accepts two 64-bit operands and produces one 64-bit result. The two operands are each recalled from an instruction specified register in either global scalar register file **211**, or M1/N1 local register file **213**. M1 unit **223** preferably performs the following instruction selected operations: 8-bit multiply operations; complex dot product operations; 32-bit bit count operations; complex conjugate multiply operations; and bit-wise Logical Operations, moves, adds and subtracts. The result may be written into an instruction specified register of global scalar register file **211**, L1/S1 local register file **212**, M1/N1 local register file **213**, or D1/D2 local register file **214**.

(58) Scalar datapath side A **115** includes N1 unit **224**. N1 unit **224** generally accepts two 64-bit operands and produces one 64-bit result. The two operands are each recalled from an instruction specified register in either global scalar register file **211** or M1/N1 local register file **213**. N1 unit **224** preferably performs the same type operations as M1 unit **223**. There may be certain double operations (called dual issued instructions) that employ both the M1 unit **223** and the N1 unit **224** together. The result may be written into an instruction specified register of global scalar register file **211**, L1/S1 local register file **212**, M1/N1 local register file **213**, or D1/D2 local register file **214**.

(59) Scalar datapath side A **115** includes D1 unit **225** and D2 unit **226**. D1 unit **225** and D2 unit **226** generally each accept two 64-bit operands and each produce one 64-bit result. D1 unit **225** and D2 unit **226** generally perform address calculations and corresponding load and store operations. D1 unit **225** is used for scalar loads and stores of 64 bits. D2 unit **226** is used for vector loads and stores of 512 bits. D1 unit **225** and D2 unit **226** preferably also perform: swapping, pack and unpack on the load and store data; 64-bit SIMD arithmetic operations; and 64-bit bit-wise logical operations. D1/D2 local register file **214** will generally store base and offset addresses used in address calculations for the corresponding loads and stores. The two operands are each recalled from an instruction specified register in either global scalar register file **211** or D1/D2 local register file **214**. The calculated result may be written into an instruction specified register of global scalar register file **211**, L1/S1 local register file **212**, M1/N1 local register file **213**, or D1/D2 local register file **214**.

(60) Vector datapath side B **116** includes L2 unit **241**. L2 unit **241** generally accepts two 512-bit operands and produces one 512-bit result. The two operands are each recalled from an instruction specified register in either global vector register file **231**, L2/S2 local register file **232** or predicate register file **234**. L2 unit **241** preferably performs instruction similar to L1 unit **221** except on wider 512-bit data. The result may be written into an instruction specified register of global vector register file **231**, L2/S2 local register file **232**, M2/N2/C local register file **233** or predicate register file **234**.

(61) Vector datapath side B **116** includes S2 unit **242**. S2 unit **242** generally accepts two 512-bit operands and produces one 512-bit result. The two operands are each recalled from an instruction specified register in either global vector register file **231**, L2/S2 local register file **232** or predicate register file **234**. S2 unit **242** preferably performs instructions similar to S1 unit **222** except on wider 512-bit data. The result may be written into an instruction specified register of global vector register file **231**, L2/S2 local register file **232**, M2/N2/C local register file **233**, or predicate register file **234**. There may be certain double operations (called dual issued instructions) that employ both L2 unit **241** and the S2 unit **242** together. The result may be written into an instruction specified register of global vector register file **231**, L2/S2 local register file **232**, or M2/N2/C local register file **233**.

(62) Vector datapath side B **116** includes M2 unit **243**. M2 unit **243** generally accepts two 512-bit operands and produces one 512-bit result. The two operands are each recalled from an instruction specified register in either global vector register file **231** or M2/N2/C local register file **233**. M2 unit **243** preferably performs instructions similar to M1 unit **223** except on wider 512-bit data. The result may be written into an instruction specified register of global vector register file **231**, L2/S2

local register file **232**, or M2/N2/C local register file **233**.

(63) Vector datapath side B **116** includes N2 unit **244**. N2 unit **244** generally accepts two 512-bit operands and produces one 512-bit result. The two operands are each recalled from an instruction specified register in either global vector register file **231** or M2/N2/C local register file **233**. N2 unit **244** preferably performs the same type operations as M2 unit **243**. There may be certain double operations (called dual issued instructions) that employ both M2 unit **243** and the N2 unit **244** together. The result may be written into an instruction specified register of global vector register file **231**, L2/S2 local register file **232**, or M2/N2/C local register file **233**.

(64) Vector datapath side B **116** includes C unit **245**. C unit **245** generally accepts two 512-bit operands and produces one 512-bit result. The two operands are each recalled from an instruction specified register in either global vector register file **231** or M2/N2/C local register file **233**. C unit **245** preferably performs: “Rake” and “Search” instructions; up to 512 2-bit PN*8-bit multiplies I/Q complex multiplies per clock cycle; 8-bit and 16-bit Sum-of-Absolute-Difference (SAD) calculations, up to 512 SADs per clock cycle; horizontal add and horizontal min/max instructions; and vector permutes instructions. C unit **245** includes also contains 4 vector control registers (CUCR0 to CUCR3) used to control certain operations of C unit **245** instructions. Control registers CUCR0 to CUCR3 are used as operands in certain C unit **245** operations. Control registers CUCR0 to CUCR3 are preferably used: in control of a general permutation instruction (VPERM); and as masks for SIMD multiple DOT product operations (DOTPM) and SIMD multiple Sum-of-Absolute-Difference (SAD) operations. Control register CUCR0 is preferably used to store the polynomials for Galois Field Multiply operations (GFMPY). Control register CUCR1 is preferably used to store the Galois field polynomial generator function.

(65) Vector datapath side B **116** includes P unit **246**. P unit **246** performs basic logic operations on registers of local predicate register file **234**. P unit **246** has direct access to read from and write to predication register file **234**. These operations include AND, ANDN, OR, XOR, NOR, BITR, NEG, SET, BITCNT, RMBD, BIT Decimate and Expand. A commonly expected use of P unit **246** includes manipulation of the SIMD vector comparison results for use in control of a further SIMD vector operation.

(66) FIG. 3 illustrates global scalar register file **211**. There are 16 independent 64-bit wide scalar registers designated A0 to A15. Each register of global scalar register file **211** can be read from or written to as 64-bits of scalar data. All scalar datapath side A **115** functional units (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226**) can read or write to global scalar register file **211**. Global scalar register file **211** may be read as 32-bits or as 64-bits and may only be written to as 64-bits. The instruction executing determines the read data size. Vector datapath side B **116** functional units (L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, C unit **245**, and P unit **246**) can read from global scalar register file **211** via crosspath **117** under restrictions that will be detailed below.

(67) FIG. 4 illustrates D1/D2 local register file **214**. There are 16 independent 64-bit wide scalar registers designated D0 to D16. Each register of D1/D2 local register file **214** can be read from or written to as 64-bits of scalar data. All scalar datapath side A **115** functional units (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226**) can write to global scalar register file **211**. Only D1 unit **225** and D2 unit **226** can read from D1/D1 local scalar register file **214**. It is expected that data stored in D1/D2 local scalar register file **214** will include base addresses and offset addresses used in address calculation.

(68) FIG. 5 illustrates L1/S1 local register file **212**. The embodiment illustrated in FIG. 5 has 8 independent 64-bit wide scalar registers designated AL0 to AL7. The preferred instruction coding (see FIG. 13) permits L1/S1 local register file **212** to include up to 16 registers. The embodiment of FIG. 5 implements only 8 registers to reduce circuit size and complexity. Each register of L1/S1 local register file **212** can be read from or written to as 64-bits of scalar data. All scalar datapath side A **115** functional units (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225**, and

D2 unit 226) can write to L1/S1 local scalar register file 212. Only L1 unit 221 and S1 unit 222 can read from L1/S1 local scalar register file 212.

(69) FIG. 6 illustrates M1/N1 local register file 213. The embodiment illustrated in FIG. 6 has 8 independent 64-bit wide scalar registers designated AM0 to AM7. The preferred instruction coding (see FIG. 13) permits M1/N1 local register file 213 to include up to 16 registers. The embodiment of FIG. 6 implements only 8 registers to reduce circuit size and complexity. Each register of M1/N1 local register file 213 can be read from or written to as 64-bits of scalar data. All scalar datapath side A 115 functional units (L1 unit 221, S1 unit 222, M1 unit 223, N1 unit 224, D1 unit 225 and D2 unit 226) can write to M1/N1 local scalar register file 213. Only M1 unit 223 and N1 unit 224 can read from M1/N1 local scalar register file 213.

(70) FIG. 7 illustrates global vector register file 231. There are 16 independent 512-bit wide scalar registers. Each register of global vector register file 231 can be read from or written to as 64-bits of scalar data designated B0 to B15. Each register of global vector register file 231 can be read from or written to as 512-bits of vector data designated VB0 to VB15. The instruction type determines the data size. All vector datapath side B 116 functional units (L2 unit 241, S2 unit 242, M2 unit 243, N2 unit 244, C unit 245, and P unit 246) can read or write to global vector register file 231. Scalar datapath side A 115 functional units (L1 unit 221, S1 unit 222, M1 unit 223, N1 unit 224, D1 unit 225, and D2 unit 226) can read from global vector register file 231 via crosspath 117 under restrictions that will be detailed below.

(71) FIG. 8 illustrates P local register file 234. There are 8 independent 64-bit wide registers designated P0 to P15. Each register of P local register file 234 can be read from or written to as 64-bits of scalar data. Vector datapath side B 116 functional units L2 unit 241, S2 unit 242, C unit 244 and P unit 246 can write to P local register file 234. Only L2 unit 241, S2 unit 242 and P unit 246 can read from P local scalar register file 234. A commonly expected use of P local register file 234 includes: writing one bit SIMD vector comparison results from L2 unit 241, S2 unit 242, or C unit 244; manipulation of the SIMD vector comparison results by P unit 246; and use of the manipulated results in control of a further SIMD vector operation.

(72) FIG. 9 illustrates L2/S2 local register file 232. The embodiment illustrated in FIG. 9 has 8 independent 512-bit wide scalar registers. The preferred instruction coding (see FIG. 13) permits L2/S2 local register file 232 to include up to 16 registers. The embodiment of FIG. 9 implements only 8 registers to reduce circuit size and complexity. Each register of L2/S2 local vector register file 232 can be read from or written to as 64-bits of scalar data designated BL0 to BL7. Each register of L2/S2 local vector register file 232 can be read from or written to as 512-bits of vector data designated VBL0 to VBL7. The instruction type determines the data size. All vector datapath side B 116 functional units (L2 unit 241, S2 unit 242, M2 unit 243, N2 unit 244, C unit 245 and P unit 246) can write to L2/S2 local vector register file 232. Only L2 unit 241 and S2 unit 242 can read from L2/S2 local vector register file 232.

(73) FIG. 10 illustrates M2/N2/C local register file 233. The embodiment illustrated in FIG. 10 has 8 independent 512-bit wide scalar registers. The preferred instruction coding (see FIG. 13) permits M2/N2/C local register file 233 to include up to 16 registers. The embodiment of FIG. 10 implements only 8 registers to reduce circuit size and complexity. Each register of M2/N2/C local vector register file 233 can be read from or written to as 64-bits of scalar data designated BM0 to BM7. Each register of M2/N2/C local vector register file 233 can be read from or written to as 512-bits of vector data designated VBM0 to VBM7. All vector datapath side B 116 functional units (L2 unit 241, S2 unit 242, M2 unit 243, N2 unit 244, C unit 245 and P unit 246) can write to M2/N2/C local vector register file 233. Only M2 unit 243, N2 unit 244 and C unit 245 can read from M2/N2/C local vector register file 233.

(74) The provision of global register files accessible by all functional units of a side and local register files accessible by only some of the functional units of a side is a design choice. This invention could be practiced employing only one type of register file corresponding to the disclosed

global register files.

(75) Crosspath **117** permits limited exchange of data between scalar datapath side A **115** and vector datapath side B **116**. During each operational cycle one 64-bit data word can be recalled from global scalar register file A **211** for use as an operand by one or more functional units of vector datapath side B **116** and one 64-bit data word can be recalled from global vector register file **231** for use as an operand by one or more functional units of scalar datapath side A **115**. Any scalar datapath side A **115** functional unit (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226**) may read a 64-bit operand from global vector register file **231**. This 64-bit operand is the least significant bits of the 512-bit data in the accessed register of global vector register file **231**. Plural scalar datapath side A **115** functional units may employ the same 64-bit crosspath data as an operand during the same operational cycle. However, only one 64-bit operand is transferred from vector datapath side B **116** to scalar datapath side A **115** in any single operational cycle. Any vector datapath side B **116** functional unit (L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, C unit **245** and P unit **246**) may read a 64-bit operand from global scalar register file **211**. If the corresponding instruction is a scalar instruction, the crosspath operand data is treated as any other 64-bit operand. If the corresponding instruction is a vector instruction, the upper 448 bits of the operand are zero filled. Plural vector datapath side B **116** functional units may employ the same 64-bit crosspath data as an operand during the same operational cycle. Only one 64-bit operand is transferred from scalar datapath side A **115** to vector datapath side B **116** in any single operational cycle.

(76) Streaming engine **125** transfers data in certain restricted circumstances. Streaming engine **125** controls two data streams. A stream consists of a sequence of elements of a particular type. Programs that operate on streams read the data sequentially, operating on each element in turn. Every stream has the following basic properties. The stream data have a well-defined beginning and ending in time. The stream data have fixed element size and type throughout the stream. The stream data have fixed sequence of elements. Thus programs cannot seek randomly within the stream. The stream data is read-only while active. Programs cannot write to a stream while simultaneously reading from it. Once a stream is opened streaming engine **125**: calculates the address; fetches the defined data type from level two unified cache (which may require cache service from a higher level memory); performs data type manipulation such as zero extension, sign extension, data element sorting/swapping such as matrix transposition; and delivers the data directly to the programmed data register file within central processing unit core **110**. Streaming engine **125** is thus useful for real-time digital filtering operations on well-behaved data. Streaming engine **125** frees these memory fetch tasks from the corresponding central processing unit core **110** enabling other processing functions.

(77) Streaming engine **125** provides the following benefits. Streaming engine **125** permits multi-dimensional memory accesses. Streaming engine **125** increases the available bandwidth to the functional units. Streaming engine **125** minimizes the number of cache miss stalls since the stream buffer bypasses level one data cache **123**. Streaming engine **125** reduces the number of scalar operations required to maintain a loop. Streaming engine **125** manages address pointers. Streaming engine **125** handles address generation automatically freeing up the address generation instruction slots and D1 unit **225** and D2 unit **226** for other computations.

(78) Central processing unit core **110** operates on an instruction pipeline. Instructions are fetched in instruction packets of fixed length further described below. All instructions require the same number of pipeline phases for fetch and decode, but require a varying number of execute phases.

(79) FIG. **11** illustrates the following pipeline phases: program fetch phase **1110**, dispatch and decode phases **1120**, and execution phases **1130**. Program fetch phase **1110** includes three stages for all instructions. Dispatch and decode phases **1120** include three stages for all instructions. Execution phase **1130** includes one to four stages dependent on the instruction.

(80) Fetch phase **1110** includes program address generation stage **1111** (PG), program access stage

1112 (PA) and program receive stage **1113 (PR)**. During program address generation stage **1111 (PG)**, the program address is generated in central processing unit core **110** and the read request is sent to the memory controller for the level one instruction cache L1I. During the program access stage **1112 (PA)** the level one instruction cache L1I processes the request, accesses the data in its memory and sends a fetch packet to the central processing unit core **110** boundary. During the program receive stage **1113 (PR)** central processing unit core **110** registers the fetch packet.

(81) Instructions are always fetched sixteen 32-bit wide slots, constituting a fetch packet, at a time. FIG. 12 illustrates 16 instructions **1201** to **1216** of a single fetch packet. Fetch packets are aligned on 512-bit (16-word) boundaries. The preferred embodiment employs a fixed 32-bit instruction length. Fixed length instructions are advantageous for several reasons. Fixed length instructions enable easy decoder alignment. A properly aligned instruction fetch can load plural instructions into parallel instruction decoders. Such a properly aligned instruction fetch can be achieved by predetermined instruction alignment when stored in memory (fetch packets aligned on 512-bit boundaries) coupled with a fixed instruction packet fetch. An aligned instruction fetch permits operation of parallel decoders on instruction-sized fetched bits. Variable length instructions require an initial step of locating each instruction boundary before they can be decoded. A fixed length instruction set generally permits more regular layout of instruction fields. This simplifies the construction of each decoder which is an advantage for a wide issue VLIW central processor.

(82) The execution of the individual instructions is partially controlled by a p bit in each instruction. This p bit is preferably bit **0** of the 32-bit wide slot. The p bit determines whether an instruction executes in parallel with a next instruction. Instructions are scanned from lower to higher address. If the p bit of an instruction is 1, then the next following instruction (higher memory address) is executed in parallel with (in the same cycle as) that instruction. If the p bit of an instruction is 0, then the next following instruction is executed in the cycle after the instruction.

(83) Central processing unit core **110** and level one instruction cache L1I **121** pipelines are decoupled from each other. Fetch packet returns from level one instruction cache L1I can take different number of clock cycles, depending on external circumstances such as whether there is a hit in level one instruction cache **121** or a hit in level two combined cache **130**. Therefore program access stage **1112 (PA)** can take several clock cycles instead of 1 clock cycle as in the other stages.

(84) The instructions executing in parallel constitute an execute packet. In the preferred embodiment an execute packet can contain up to sixteen instructions. No two instructions in an execute packet may use the same functional unit. A slot is one of five types: 1) a self-contained instruction executed on one of the functional units of central processing unit core **110** (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225**, D2 unit **226**, L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, C unit **245** and P unit **246**); 2) a unitless instruction such as a NOP (no operation) instruction or multiple NOP instruction; 3) a branch instruction; 4) a constant field extension; and 5) a conditional code extension. Some of these slot types will be further explained below.

(85) Dispatch and decode phases **1110** include instruction dispatch to appropriate execution unit stage **1121 (DS)**, instruction pre-decode stage **1122 (DC1)**, and instruction decode, operand reads stage **1123 (DC2)**. During instruction dispatch to appropriate execution unit stage **1121 (DS)**, the fetch packets are split into execute packets and assigned to the appropriate functional units. During the instruction pre-decode stage **1122 (DC1)**, the source registers, destination registers and associated paths are decoded for the execution of the instructions in the functional units. During the instruction decode, operand reads stage **1123 (DC2)**, more detailed unit decodes are done, as well as reading operands from the register files.

(86) Execution phases **1130** includes execution stages **1131** to **1135 (E1 to E5)**. Different types of instructions require different numbers of these stages to complete their execution. These stages of the pipeline play an important role in understanding the device state at central processing unit core **110** cycle boundaries.

(87) During execute **1** stage **1131 (E1)** the conditions for the instructions are evaluated and

operands are operated on. As illustrated in FIG. 11, execute 1 stage **1131** may receive operands from a stream buffer **1141** and one of the register files shown schematically as **1142**. For load and store instructions, address generation is performed and address modifications are written to a register file. For branch instructions, branch fetch packet in PG phase **1111** is affected. As illustrated in FIG. 11, load and store instructions access memory here shown schematically as memory **1151**. For single-cycle instructions, results are written to a destination register file. This assumes that any conditions for the instructions are evaluated as true. If a condition is evaluated as false, the instruction does not write any results or have any pipeline operation after execute 1 stage **1131**.

(88) During execute 2 stage **1132** (E2) load instructions send the address to memory. Store instructions send the address and data to memory. Single-cycle instructions that saturate results set the SAT bit in the control status register (CSR) if saturation occurs. For 2-cycle instructions, results are written to a destination register file.

(89) During execute 3 stage **1133** (E3) data memory accesses are performed. Any multiply instructions that saturate results set the SAT bit in the control status register (CSR) if saturation occurs. For 3-cycle instructions, results are written to a destination register file.

(90) During execute 4 stage **1134** (E4) load instructions bring data to the central processing unit core **110** boundary. For 4-cycle instructions, results are written to a destination register file.

(91) During execute 5 stage **1135** (E5) load instructions write data into a register. This is illustrated schematically in FIG. 11 with input from memory **1151** to execute 5 stage **1135**.

(92) FIG. 13 illustrates an example of the instruction coding **1300** of functional unit instructions used by this invention. Those skilled in the art would realize that other instruction codings are feasible and within the scope of this invention. Each instruction consists of 32 bits and controls the operation of one of the individually controllable functional units (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225**, D2 unit **226**, L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, C unit **245** and P unit **246**). The bit fields are defined as follows.

(93) The creg field **1301** (bits **29** to **31**) and the z bit **1302** (bit **28**) are optional fields used in conditional instructions. These bits are used for conditional instructions to identify the predicate register and the condition. The z bit **1302** (bit **28**) indicates whether the predication is based upon zero or not zero in the predicate register. If z=1, the test is for equality with zero. If z=0, the test is for nonzero. The case of creg=0 and z=0 is treated as always true to allow unconditional instruction execution. The creg field **1301** and the z field **1302** are encoded in the instruction as shown in Table 1.

(94) TABLE-US-00001 TABLE 1 Conditional creg z Register 31 30 29 28 Unconditional 0 0 0 0
Reserved 0 0 0 1 A0 0 0 1 z A1 0 1 0 z A2 0 1 1 z A3 1 0 0 z A4 1 0 1 z A5 1 1 0 z Reserved 1 1 x x

(95) Execution of a conditional instruction is conditional upon the value stored in the specified data register. This data register is in the global scalar register file **211** for all functional units. Note that “z” in the z bit column refers to the zero/not zero comparison selection noted above and “x” is a don't care state. This coding can only specify a subset of the **16** global registers as predicate registers. This selection was made to preserve bits in the instruction coding. Note that unconditional instructions do not have these optional bits. For unconditional instructions these bits in fields **1301** and **1302** (**28** to **31**) are preferably used as additional opcode bits.

(96) The dst field **1303** (bits **23** to **27**) specifies a register in a corresponding register file as the destination of the instruction results.

(97) The src2/cst field **1304** (bits **18** to **22**) has several meanings depending on the instruction opcode field (bits **3** to **12** for all instructions and additionally bits **28** to **31** for unconditional instructions). The first meaning specifies a register of a corresponding register file as the second operand. The second meaning is an immediate constant. Depending on the instruction type, this is treated as an unsigned integer and zero extended to a specified data length or is treated as a signed integer and sign extended to the specified data length.

(98) The src1 field **1305** (bits **13** to **17**) specifies a register in a corresponding register file as the first source operand.

(99) The opcode field **1306** (bits **3** to **12**) for all instructions (and additionally bits **28** to **31** for unconditional instructions) specifies the type of instruction and designates appropriate instruction options. This includes unambiguous designation of the functional unit used and operation performed. A detailed explanation of the opcode is beyond the scope of this invention except for the instruction options detailed below.

(100) The e bit **1307** (bit **2**) is only used for immediate constant instructions where the constant may be extended. If e=1, then the immediate constant is extended in a manner detailed below. If e=0, then the immediate constant is not extended. In that case the immediate constant is specified by the src2/cst field **1304** (bits **18** to **22**). Note that this e bit **1307** is used for only some instructions. Accordingly, with proper coding this e bit **1307** may be omitted from instructions which do not need it and this bit used as an additional opcode bit.

(101) The s bit **1308** (bit **1**) designates scalar datapath side A **115** or vector datapath side B **116**. If s=0, then scalar datapath side A **115** is selected. This limits the functional unit to L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226** and the corresponding register files illustrated in FIG. 2. Similarly, s=1 selects vector datapath side B **116** limiting the functional unit to L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244**, P unit **246** and the corresponding register file illustrated in FIG. 2.

(102) The p bit **1309** (bit **0**) marks the execute packets. The p-bit determines whether the instruction executes in parallel with the following instruction. The p-bits are scanned from lower to higher address. If p=1 for the current instruction, then the next instruction executes in parallel with the current instruction. If p=0 for the current instruction, then the next instruction executes in the cycle after the current instruction. All instructions executing in parallel constitute an execute packet. An execute packet can contain up to twelve instructions. Each instruction in an execute packet must use a different functional unit.

(103) There are two different condition code extension slots. Each execute packet can contain one each of these unique 32-bit condition code extension slots which contains the 4-bit creg/z fields for the instructions in the same execute packet. FIG. 14 illustrates the coding for condition code extension slot **0** and FIG. 15 illustrates the coding for condition code extension slot **1**.

(104) FIG. 14 illustrates the coding for condition code extension slot **0** having 32 bits. Field **1401** (bits **28** to **31**) specify **4** creg/z bits assigned to the L1 unit **221** instruction in the same execute packet. Field **1402** (bits **27** to **24**) specify **4** creg/z bits assigned to the L2 unit **241** instruction in the same execute packet. Field **1403** (bits **19** to **23**) specify **4** creg/z bits assigned to the S1 unit **222** instruction in the same execute packet. Field **1404** (bits **16** to **19**) specify **4** creg/z bits assigned to the S2 unit **242** instruction in the same execute packet. Field **1405** (bits **12** to **15**) specify **4** creg/z bits assigned to the D1 unit **225** instruction in the same execute packet. Field **1406** (bits **8** to **11**) specify **4** creg/z bits assigned to the D2 unit **226** instruction in the same execute packet. Field **1407** (bits **6** and **7**) is unused/reserved. Field **1408** (bits **0** to **5**) are coded a set of unique bits (CCEX0) to identify the condition code extension slot **0**. Once this unique ID of condition code extension slot **0** is detected, the corresponding creg/z bits are employed to control conditional execution of any L1 unit **221**, L2 unit **241**, S1 unit **222**, S2 unit **242**, D1 unit **225** and D2 unit **226** instruction in the same execution packet. These creg/z bits are interpreted as shown in Table 1. If the corresponding instruction is conditional (includes creg/z bits) the corresponding bits in the condition code extension slot **0** override the condition code bits in the instruction. Note that no execution packet can have more than one instruction directed to a particular execution unit. No execute packet of instructions can contain more than one condition code extension slot **0**. Thus the mapping of creg/z bits to functional unit instruction is unambiguous. Setting the creg/z bits equal to "0000" makes the instruction unconditional. Thus a properly coded condition code extension slot **0** can make some corresponding instructions conditional and some unconditional.

(105) FIG. 15 illustrates the coding for condition code extension slot 1 having 32 bits. Field 1501 (bits 28 to 31) specify 4 creg/z bits assigned to the M1 unit 223 instruction in the same execute packet. Field 1502 (bits 27 to 24) specify 4 creg/z bits assigned to the M2 unit 243 instruction in the same execute packet. Field 1503 (bits 19 to 23) specify 4 creg/z bits assigned to the C unit 245 instruction in the same execute packet. Field 1504 (bits 16 to 19) specify 4 creg/z bits assigned to the N1 unit 224 instruction in the same execute packet. Field 1505 (bits 12 to 15) specify 4 creg/z bits assigned to the N2 unit 244 instruction in the same execute packet. Field 1506 (bits 6 to 11) is unused/reserved. Field 1507 (bits 0 to 5) are coded a set of unique bits (CCEX1) to identify the condition code extension slot 1. Once this unique ID of condition code extension slot 1 is detected, the corresponding creg/z bits are employed to control conditional execution of any M1 unit 223, M2 unit 243, C unit 245, N1 unit 224 and N2 unit 244 instruction in the same execution packet. These creg/z bits are interpreted as shown in Table 1. If the corresponding instruction is conditional (includes creg/z bits) the corresponding bits in the condition code extension slot 1 override the condition code bits in the instruction. Note that no execution packet can have more than one instruction directed to a particular execution unit. No execute packet of instructions can contain more than one condition code extension slot 1. Thus the mapping of creg/z bits to functional unit instruction is unambiguous. Setting the creg/z bits equal to "0000" makes the instruction unconditional. Thus a properly coded condition code extension slot 1 can make some instructions conditional and some unconditional.

(106) It is feasible for both condition code extension slot 0 and condition code extension slot 1 to include a p bit to define an execute packet as described above in conjunction with FIG. 13. In the preferred embodiment, as illustrated in FIGS. 14 and 15, code extension slot 0 and condition code extension slot 1 preferably have bit 0 (p bit) always encoded as 1. Thus neither condition code extension slot 0 nor condition code extension slot 1 can be in the last instruction slot of an execute packet.

(107) There are two different constant extension slots. Each execute packet can contain one each of these unique 32-bit constant extension slots which contains 27 bits to be concatenated as high order bits with the 5-bit constant field 1305 to form a 32-bit constant. As noted in the instruction coding description above only some instructions define the src2/cst field 1304 as a constant rather than a source register identifier. At least some of those instructions may employ a constant extension slot to extend this constant to 32 bits.

(108) FIG. 16 illustrates the fields of constant extension slot 0. Each execute packet may include one instance of constant extension slot 0 and one instance of constant extension slot 1. FIG. 16 illustrates that constant extension slot 0 1600 includes two fields. Field 1601 (bits 5 to 31) constitute the most significant 27 bits of an extended 32-bit constant including the target instruction scr2/cst field 1304 as the five least significant bits. Field 1602 (bits 0 to 4) are coded a set of unique bits (CSTX0) to identify the constant extension slot 0. In the preferred embodiment constant extension slot 0 1600 can only be used to extend the constant of one of an L1 unit 221 instruction, data in a D1 unit 225 instruction, an S2 unit 242 instruction, an offset in a D2 unit 226 instruction, an M2 unit 243 instruction, an N2 unit 244 instruction, a branch instruction, or a C unit 245 instruction in the same execute packet. Constant extension slot 1 is similar to constant extension slot 0 except that bits 0 to 4 are coded a set of unique bits (CSTX1) to identify the constant extension slot 1. In the preferred embodiment constant extension slot 1 can only be used to extend the constant of one of an L2 unit 241 instruction, data in a D2 unit 226 instruction, an S1 unit 222 instruction, an offset in a D1 unit 225 instruction, an M1 unit 223 instruction or an N1 unit 224 instruction in the same execute packet.

(109) Constant extension slot 0 and constant extension slot 1 are used as follows. The target instruction must be of the type permitting constant specification. As known in the art this is implemented by replacing one input operand register specification field with the least significant bits of the constant as described above with respect to scr2/cst field 1304. Instruction decoder 113

determines this case, known as an immediate field, from the instruction opcode bits. The target instruction also includes one constant extension bit (e bit **1307**) dedicated to signaling whether the specified constant is not extended (preferably constant extension bit=0) or the constant is extended (preferably constant extension bit=1). If instruction decoder **113** detects a constant extension slot **0** or a constant extension slot **1**, it further checks the other instructions within that execute packet for an instruction corresponding to the detected constant extension slot. A constant extension is made only if one corresponding instruction has a constant extension bit (e bit **1307**) equal to 1.

(110) FIG. **17** is a partial block diagram **1700** illustrating constant extension. FIG. **17** assumes that instruction decoder **113** detects a constant extension slot and a corresponding instruction in the same execute packet. Instruction decoder **113** supplies the 27 extension bits from the constant extension slot (bit field **1601**) and the 5 constant bits (bit field **1305**) from the corresponding instruction to concatenator **1701**. Concatenator **1701** forms a single 32-bit word from these two parts. In the preferred embodiment the 27 extension bits from the constant extension slot (bit field **1601**) are the most significant bits and the 5 constant bits (bit field **1305**) are the least significant bits. This combined 32-bit word is supplied to one input of multiplexer **1702**. The 5 constant bits from the corresponding instruction field **1305** supply a second input to multiplexer **1702**. Selection of multiplexer **1702** is controlled by the status of the constant extension bit. If the constant extension bit (e bit **1307**) is 1 (extended), multiplexer **1702** selects the concatenated 32-bit input. If the constant extension bit is 0 (not extended), multiplexer **1702** selects the 5 constant bits from the corresponding instruction field **1305**. Multiplexer **1702** supplies this output to an input of sign extension unit **1703**.

(111) Sign extension unit **1703** forms the final operand value from the input from multiplexer **1703**. Sign extension unit **1703** receives control inputs Scalar/Vector and Data Size. The Scalar/Vector input indicates whether the corresponding instruction is a scalar instruction or a vector instruction. The functional units of data path side A **115** (L1 unit **221**, S1 unit **222**, M1 unit **223**, N1 unit **224**, D1 unit **225** and D2 unit **226**) can only perform scalar instructions. Any instruction directed to one of these functional units is a scalar instruction. Data path side B functional units L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244** and C unit **245** may perform scalar instructions or vector instructions. Instruction decoder **113** determines whether the instruction is a scalar instruction or a vector instruction from the opcode bits. P unit **246** may only perform scalar instructions. The Data Size may be 8 bits (byte B), 16 bits (half-word H), 32 bits (word W) or 64 bits (double word D).

(112) Table 2 lists the operation of sign extension unit **1703** for the various options.

(113) TABLE-US-00002

TABLE 2	Instruction	Operand	Constant	Type	Size	Length	Action
Scalar	B/H/W/D	5 bits	Sign extend	to 64 bits	Scalar	B/H/W/D	32 bits
Vector	B/H/W/D	5 bits	Sign extend	to operand size and replicate across whole	Vector	B/H/W	32 bits
Vector	D	32 bits	Replicate	32-bit constant across each 32-bit (W) lane	Vector	D	32 bits
Vector	D	32 bits	Sign extend	to 64 bits and replicate across each 64-bit (D) lane	Vector	D	32 bits

(114) It is feasible for both constant extension slot **0** and constant extension slot **1** to include a p bit to define an execute packet as described above in conjunction with FIG. **13**. In the preferred embodiment, as in the case of the condition code extension slots, constant extension slot **0** and constant extension slot **1** preferably have bit **0** (p bit) always encoded as **1**. Thus neither constant extension slot **0** nor constant extension slot **1** can be in the last instruction slot of an execute packet.

(115) It is technically feasible for an execute packet to include a constant extension slot **0** or **1** and more than one corresponding instruction marked constant extended (e bit=1). For constant extension slot **0** this would mean more than one of an L1 unit **221** instruction, data in a D1 unit **225** instruction, an S2 unit **242** instruction, an offset in a D2 unit **226** instruction, an M2 unit **243** instruction or an N2 unit **244** instruction in an execute packet have an e bit of 1. For constant extension slot **1** this would mean more than one of an L2 unit **241** instruction, data in a D2 unit **226** instruction, an S1 unit **222** instruction, an offset in a D1 unit **225** instruction, an M1 unit **223** instruction or an N1 unit **224** instruction in an execute packet have an e bit of 1. Supplying the

same constant extension to more than one instruction is not expected to be a useful function. Accordingly, in one embodiment instruction decoder **113** may determine this case an invalid operation and not supported. Alternately, this combination may be supported with extension bits of the constant extension slot applied to each corresponding functional unit instruction marked constant extended.

(116) Special vector predicate instructions use registers in predicate register file **234** to control vector operations. In the current embodiment all these SIMD vector predicate instructions operate on selected data sizes. The data sizes may include byte (8 bit) data, half word (16 bit) data, word (32 bit) data, double word (64 bit) data, quad word (128 bit) data and half vector (256 bit) data. Each bit of the predicate register controls whether a SIMD operation is performed upon the corresponding byte of data. The operations of P unit **245** permit a variety of compound vector SIMD operations based upon more than one vector comparison. For example a range determination can be made using two comparisons. A candidate vector is compared with a first vector reference having the minimum of the range packed within a first data register. A second comparison of the candidate vector is made with a second reference vector having the maximum of the range packed within a second data register. Logical combinations of the two resulting predicate registers would permit a vector conditional operation to determine whether each data part of the candidate vector is within range or out of range.

(117) L1 unit **221**, S1 unit **222**, L2 unit **241**, S2 unit **242** and C unit **245** often operate in a single instruction multiple data (SIMD) mode. In this SIMD mode the same instruction is applied to packed data from the two operands. Each operand holds plural data elements disposed in predetermined slots. SIMD operation is enabled by carry control at the data boundaries. Such carry control enables operations on varying data widths.

(118) FIG. **18** illustrates the carry control. AND gate **1801** receives the carry output of bit N within the operand wide arithmetic logic unit (64 bits for scalar datapath side A **115** functional units and 512 bits for vector datapath side B **116** functional units). AND gate **1801** also receives a carry control signal which will be further explained below. The output of AND gate **1801** is supplied to the carry input of bit N+1 of the operand wide arithmetic logic unit. AND gates such as AND gate **1801** are disposed between every pair of bits at a possible data boundary. For example, for 8-bit data such an AND gate will be between bits **7** and **8**, bits **15** and **16**, bits **23** and **24**, etc. Each such AND gate receives a corresponding carry control signal. If the data size is the minimum, then each carry control signal is 0, effectively blocking carry transmission between the adjacent bits. The corresponding carry control signal is 1 if the selected data size requires both arithmetic logic unit sections. Table 3 below shows example carry control signals for the case of a 512 bit wide operand such as used by vector datapath side B **116** functional units which may be divided into sections of 8 bits, 16 bits, 32 bits, 64 bits, 128 bits or 256 bits. In Table 3 the upper 32 bits control the upper bits (bits **128** to **511**) carries and the lower 32 bits control the lower bits (bits **0** to **127**) carries. No control of the carry output of the most significant bit is needed, thus only 63 carry control signals are required.

(119) TABLE-US-00003 TABLE 3 Data Size Carry Control Signals

8 bits (B)	-000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000 0000
16 bits (H)	-101 0101 0101 0101 0101 0101 0101 0101 0101 0101 0101 0101 0101 0101
32 bits (W)	-111 0111 0111 0111 0111 0111 0111 0111 0111 0111 0111 0111 0111 0111
64 bits (D)	-111 1111 0111 1111 1111 0111 1111 0111 1111 0111 1111 0111 1111 1111
128 bits	-111 1111 1111 0111 1111 1111 0111 1111 1111 0111 1111 1111 1111 1111
256 bits	-111 1111 1111 1111 0111 1111 1111 1111 1111 1111 1111 1111 1111 1111

(120) It is typical in the art to operate on data sizes that are integral powers of 2 ($2^{\text{sup.N}}$). However, this carry control technique is not limited to integral powers of 2. One skilled in the art would understand how to apply this technique to other data sizes and other operand widths.

(121) FIG. **19** illustrates a conceptual view of the streaming engines of this invention. FIG. **19**

illustrates the process of a single stream. Streaming engine **1900** includes stream address generator **1901**. Stream address generator **1901** sequentially generates addresses of the elements of the stream and supplies these element addresses to system memory **1910**. Memory **1910** recalls data stored at the element addresses (data elements) and supplies these data elements to data first-in-first-out (FIFO) memory **1902**. Data FIFO **1902** provides buffering between memory **1910** and CPU **1920**. Data formatter **1903** receives the data elements from data FIFO memory **1902** and provides data formatting according to the stream definition. This process will be described below. Streaming engine **1900** supplies the formatted data elements from data formatter **1903** to the CPU **1920**. The program on CPU **1920** consumes the data and generates an output.

(122) Stream elements typically reside in normal memory. The memory itself imposes no particular structure upon the stream. Programs define streams and therefore impose structure, by specifying the following stream attributes: address of the first element of the stream; size and type of the elements in the stream; formatting for data in the stream; and the address sequence associated with the stream.

(123) The streaming engine defines an address sequence for elements of the stream in terms of a pointer walking through memory. A multiple-level nested loop controls the path the pointer takes. An iteration count for a loop level indicates the number of times that level repeats. A dimension gives the distance between pointer positions of that loop level.

(124) In a basic forward stream the innermost loop always consumes physically contiguous elements from memory. The implicit dimension of this innermost loop is 1 element. The pointer itself moves from element to element in consecutive, increasing order. In each level outside the inner loop, that loop moves the pointer to a new location based on the size of that loop level's dimension.

(125) This form of addressing allows programs to specify regular paths through memory in a small number of parameters. Table 4 lists the addressing parameters of a basic stream.

(126) TABLE-US-00004 TABLE 4 Parameter Definition
ELEM_BYTES Size of each element in bytes
ICNT0 Number of iterations for the innermost loop level 0. At loop level 0 all elements are physically contiguous
DIM0 is **ELEM_BYTES**
ICNT1 Number of iterations for loop level 1
DIM1 Number of bytes between the starting points for consecutive iterations of loop level 1
ICNT2 Number of iterations for loop level 2
DIM2 Number of bytes between the starting points for consecutive iterations of loop level 2
ICNT3 Number of iterations for loop level 3
DIM3 Number of bytes between the starting points for consecutive iterations of loop level 3
ICNT4 Number of iterations for loop level 4
DIM4 Number of bytes between the starting points for consecutive iterations of loop level 4
ICNT5 Number of iterations for loop level 5
DIM5 Number of bytes between the starting points for consecutive iterations of loop level 5

(127) The definition above maps consecutive elements of the stream to increasing addresses in memory. This works well for most algorithms but not all. Some algorithms are better served by reading elements in decreasing memory addresses, reverse stream addressing. For example, a discrete convolution computes vector dot-products, as per the formula:

$$(128) (f, g)[t] = \sum_{x=-\infty}^{\infty} f[x]g[t-x]$$

(129) In most DSP code, $f[\]$ and $g[\]$ represent arrays in memory. For each output, the algorithm reads $f[\]$ in the forward direction, but reads $g[\]$ in the reverse direction. Practical filters limit the range of indices for $[x]$ and $[t-x]$ to a finite number elements. To support this pattern, the streaming engine supports reading elements in decreasing address order.

(130) Matrix multiplication presents a unique problem to the streaming engine. Each element in the matrix product is a vector dot product between a row from the first matrix and a column from the second. Programs typically store matrices all in row-major or column-major order. Row-major order stores all the elements of a single row contiguously in memory. Column-major order stores all elements of a single column contiguously in memory. Matrices typically get stored in the same

order as the default array order for the language. As a result, only one of the two matrices in a matrix multiplication map on to the streaming engine's 2-dimensional stream definition. In a typical example a first index steps through columns on array first array but rows on second array. This problem is not unique to the streaming engine. Matrix multiplication's access pattern fits poorly with most general-purpose memory hierarchies. Some software libraries transposed one of the two matrices, so that both get accessed row-wise (or column-wise) during multiplication. The streaming engine supports implicit matrix transposition with transposed streams. Transposed streams avoid the cost of explicitly transforming the data in memory. Instead of accessing data in strictly consecutive-element order, the streaming engine effectively interchanges the inner two loop dimensions in its traversal order, fetching elements along the second dimension into contiguous vector lanes.

(131) This algorithm works, but is impractical to implement for small element sizes. Some algorithms work on matrix tiles which are multiple columns and rows together. Therefore, the streaming engine defines a separate transposition granularity. The hardware imposes a minimum granularity. The transpose granularity must also be at least as large as the element size.

Transposition granularity causes the streaming engine to fetch one or more consecutive elements from dimension 0 before moving along dimension 1. When the granularity equals the element size, this results in fetching a single column from a row-major array. Otherwise, the granularity specifies fetching 2, 4 or more columns at a time from a row-major array. This is also applicable for column-major layout by exchanging row and column in the description. A parameter GRANULE indicates the transposition granularity in bytes.

(132) Another common matrix multiplication technique exchanges the innermost two loops of the matrix multiply. The resulting inner loop no longer reads down the column of one matrix while reading across the row of another. For example the algorithm may hoist one term outside the inner loop, replacing it with the scalar value. On a vector machine, the innermost loop can be implemented very efficiently with a single scalar-by-vector multiply followed by a vector add. The central processing unit core **110** of this invention lacks a scalar-by-vector multiply. Programs must instead duplicate the scalar value across the length of the vector and use a vector-by-vector multiply. The streaming engine of this invention directly supports this and related use models with an element duplication mode. In this mode, the streaming engine reads a granule smaller than the full vector size and replicates that granule to fill the next vector output.

(133) The streaming engine treats each complex number as a single element with two sub-elements that give the real and imaginary (rectangular) or magnitude and angle (polar) portions of the complex number. Not all programs or peripherals agree what order these sub-elements should appear in memory. Therefore, the streaming engine offers the ability to swap the two sub-elements of a complex number with no cost. This feature swaps the halves of an element without interpreting the contents of the element and can be used to swap pairs of sub-elements of any type, not just complex numbers.

(134) Algorithms generally prefer to work at high precision, but high precision values require more storage and bandwidth than lower precision values. Commonly, programs will store data in memory at low precision, promote those values to a higher precision for calculation and then demote the values to lower precision for storage. The streaming engine supports this directly by allowing algorithms to specify one level of type promotion. In the preferred embodiment of this invention every sub-element may be promoted to the next larger type size with either sign or zero extension for integer types. It is also feasible that the streaming engine may support floating point promotion, promoting 16-bit and 32-bit floating point values to 32-bit and 64-bit formats, respectively.

(135) The streaming engine defines a stream as a discrete sequence of data elements, the central processing unit core **110** consumes data elements packed contiguously in vectors. Vectors resemble streams in as much as they contain multiple homogeneous elements with some implicit sequence.

Because the streaming engine reads streams, but the central processing unit core **110** consumes vectors, the streaming engine must map streams onto vectors in a consistent way.

(136) Vectors consist of equal-sized lanes, each lane containing a sub-element. The central processing unit core **110** designates the rightmost lane of the vector as lane **0**, regardless of device's current endian mode. Lane numbers increase right-to-left. The actual number of lanes within a vector varies depending on the length of the vector and the data size of the sub-element.

(137) FIG. **20** illustrates a first example of lane allocation in a vector. Vector **2000** is divided into 8 64-bit lanes ($8 \times 64 \text{ bits} = 512 \text{ bits}$ the vector length). Lane **0** includes bits **0** to **63**; line **1** includes bits **64** to **127**; lane **2** includes bits **128** to **191**; lane **3** includes bits **192** to **255**, lane **4** includes bits **256** to **319**, lane **5** includes bits **320** to **383**, lane **6** includes bits **384** to **447** and lane **7** includes bits **448** to **511**.

(138) FIG. **21** illustrates a second example of lane allocation in a vector. Vector **2100** is divided into 16 32-bit lanes ($16 \times 32 \text{ bits} = 512 \text{ bits}$ the vector length). Lane **0** includes bits **0** to **31**; line **1** includes bits **32** to **63**; lane **2** includes bits **64** to **95**; lane **3** includes bits **96** to **127**; lane **4** includes bits **128** to **159**; lane **5** includes bits **160** to **191**; lane **6** includes bits **192** to **223**; lane **7** includes bits **224** to **255**; lane **8** includes bits **256** to **287**; line **9** occupied bits **288** to **319**; lane **10** includes bits **320** to **351**; lane **11** includes bits **352** to **383**; lane **12** includes bits **384** to **415**; lane **13** includes bits **416** to **447**; lane **14** includes bits **448** to **479**; and lane **15** includes bits **480** to **511**.

(139) The streaming engine maps the innermost stream dimension directly to vector lanes. It maps earlier elements within that dimension to lower lane numbers and later elements to higher lane numbers. This is true regardless of whether this particular stream advances in increasing or decreasing address order. Whatever order the stream defines, the streaming engine deposits elements in vectors in increasing-lane order. For non-complex data, it places the first element in lane **0** of the first vector central processing unit core **110** fetches, the second in lane **1**, and so on. For complex data, the streaming engine places the first element in lanes **0** and **1**, second in lanes **2** and **3**, and so on. Sub-elements within an element retain the same relative ordering regardless of the stream direction. For non-swapped complex elements, this places the sub-elements with the lower address of each pair in the even numbered lanes, and the sub-elements with the higher address of each pair in the odd numbered lanes. Swapped complex elements reverse this mapping.

(140) The streaming engine fills each vector central processing unit core **110** fetches with as many elements as it can from the innermost stream dimension. If the innermost dimension is not a multiple of the vector length, the streaming engine pads that dimension out to a multiple of the vector length with zeros. Thus for higher-dimension streams, the first element from each iteration of an outer dimension arrives in lane **0** of a vector. The streaming engine always maps the innermost dimension to consecutive lanes in a vector. For transposed streams, the innermost dimension consists of groups of sub-elements along dimension 1, not dimension 0, as transposition exchanges these two dimensions.

(141) Two dimensional streams exhibit great variety as compared to one dimensional streams. A basic two dimensional stream extracts a smaller rectangle from a larger rectangle. A transposed 2-D stream reads a rectangle column-wise instead of row-wise. A looping stream, where the second dimension overlaps first executes a finite impulse response (FIR) filter taps which loops repeatedly or FIR filter samples which provide a sliding window of input samples.

(142) FIG. **22** illustrates a basic two dimensional stream. The inner two dimensions, represented by ELEM_BYTES, ICNT0, DIM1 and ICNT1 give sufficient flexibility to describe extracting a smaller rectangle **2220** having dimensions **2221** and **2222** from a larger rectangle **2210** having dimensions **2211** and **2212**. In this example rectangle **2220** is a 9 by 13 rectangle of 64-bit values and rectangle **2210** is a larger 11 by 19 rectangle. The following stream parameters define this stream: ICNT0=9 ELEM_BYTES=8 ICNT1=13 DIM1=88 (11 times 8)

(143) Thus the iteration count in the 0 dimension **2221** is 9. The iteration count in the 1 direction **2222** is 13. Note that the ELEM_BYTES only scales the innermost dimension. The first dimension

has ICNT0 elements of size ELEM_BYTES. The stream address generator does not scale the outer dimensions. Therefore, DIM1=88, which is 11 elements scaled by 8 bytes per element.

(144) FIG. 23 illustrates the order of elements within this example stream. The streaming engine fetches elements for the stream in the order illustrated in order 2300. The first 9 elements come from the first row of rectangle 2220, left-to-right in hops 1 to 8. The 10th through 24th elements comes from the second row, and so on. When the stream moves from the 9th element to the 10th element (hop 9 in FIG. 23), the streaming engine computes the new location based on the pointer's position at the start of the inner loop, not where the pointer ended up at the end of the first dimension. This makes DIM1 independent of ELEM_BYTES and ICNT0. DIM1 always represents the distance between the first bytes of each consecutive row.

(145) Transposed streams access along dimension 1 before dimension 0. The following examples illustrate a couple transposed streams, varying the transposition granularity. FIG. 24 illustrates extracting a smaller rectangle 2420 (12×8) having dimensions 2421 and 2422 from a larger rectangle 2410 (14×13) having dimensions 2411 and 2412. In FIG. 24 ELEM_BYTES equals 2.

(146) FIG. 25 illustrates how the streaming engine would fetch the stream of this example with a transposition granularity of 4 bytes. Fetch pattern 2500 fetches pairs of elements from each row (because the granularity of 4 is twice the ELEM_BYTES of 2), but otherwise moves down the columns. Once it reaches the bottom of a pair of columns, it repeats this pattern with the next pair of columns.

(147) FIG. 26 illustrates how the streaming engine would fetch the stream of this example with a transposition granularity of 8 bytes. The overall structure remains the same. The streaming engine fetches 4 elements from each row (because the granularity of 8 is four times the ELEM_BYTES of 2) before moving to the next row in the column as shown in fetch pattern 2600.

(148) The streams examined so far read each element from memory exactly once. A stream can read a given element from memory multiple times, in effect looping over a piece of memory. FIR filters exhibit two common looping patterns. FIRs re-read the same filter taps for each output. FIRs also read input samples from a sliding window. Two consecutive outputs will need inputs from two overlapping windows.

(149) FIG. 27 illustrates the details of streaming engine 2700. Streaming engine 2700 contains three major sections: Stream 0 2710; Stream 1 2720; and Shared L2 Interfaces 2730. Stream 0 2710 and Stream 1 2720 both contain identical hardware that operates in parallel. Stream 0 2710 and Stream 1 2720 both share L2 interfaces 2730. Each stream 2710 and 2720 provides central processing unit core 110 with up to 512 bits/cycle, every cycle. The streaming engine architecture enables this through its dedicated stream paths and shared dual L2 interfaces.

(150) Each streaming engine 2700 includes a dedicated 6-dimensional stream address generator 2711/2721 that can each generate one new non-aligned request per cycle. Address generators 2711/2721 output 512-bit aligned addresses that overlap the elements in the sequence defined by the stream parameters. This will be further described below.

(151) Each address generator 2711/2721 connects to a dedicated micro table look-aside buffer (μTLB) 2712/2722. The μTLB 2712/2722 converts a single 48-bit virtual address to a 44-bit physical address each cycle. Each μTLB 2712/2722 has 8 entries, covering a minimum of 32 kB with 4 kB pages or a maximum of 16 MB with 2 MB pages. Each address generator 2711/2721 generates 2 addresses per cycle. The μTLB 2712/2722 only translates 1 address per cycle. To maintain throughput, streaming engine 2700 takes advantage of the fact that most stream references will be within the same 4 kB page. Thus the address translation does not modify bits 0 to 11 of the address. If aout0 and aout1 line in the same 4 kB page (aout0[47:12] are the same as aout1[47:12]), then the μTLB 2712/2722 only translates aout0 and reuses the translation for the upper bits of both addresses.

(152) Translated addresses are queued in command queue 2713/2723. These addresses are aligned with information from the corresponding Storage Allocation and Tracking block 2714/2724.

Streaming engine **2700** does not explicitly manage μ TLB **2712/2722**. The system memory management unit (MMU) invalidates μ TLBs as necessary during context switches.

(153) Storage Allocation and Tracking **2714/2724** manages the stream's internal storage, discovering data reuse and tracking the lifetime of each piece of data. This will be further described below.

(154) Reference queue **2715/2725** stores the sequence of references generated by the corresponding address generator **2711/2721**. This information drives the data formatting network so that it can present data to central processing unit core **110** in the correct order. Each entry in reference queue **2715/2725** contains the information necessary to read data out of the data store and align it for central processing unit core **110**. Reference queue **2715/2725** maintains the following information listed in Table 5 in each slot:

(155) TABLE-US-00005 TABLE 5 Data Slot Low Slot number for the lower half of data associated with aout0 Data Slot High Slot number for the upper half of data associated with aout1 Rotation Number of bytes to rotate data to align next element with lane 0 Length Number of valid bytes in this reference

(156) Storage allocation and tracking **2714/2724** inserts references in reference queue **2715/2725** as address generator **2711/2721** generates new addresses. Storage allocation and tracking **2714/2724** removes references from reference queue **2715/2725** when the data becomes available and there is room in the stream holding registers. As storage allocation and tracking **2714/2724** removes slot references from reference queue **2715/2725** and formats data, it checks whether the references represent the last reference to the corresponding slots. Storage allocation and tracking **2714/2724** compares reference queue **2715/2725** removal pointer against the slot's recorded Last Reference. If they match, then storage allocation and tracking **2714/2724** marks the slot inactive once it's done with the data.

(157) Streaming engine **2700** has data storage **2716/2726** for an arbitrary number of elements. Deep buffering allows the streaming engine to fetch far ahead in the stream, hiding memory system latency. The right amount of buffering might vary from product generation to generation. In the current preferred embodiment streaming engine **2700** dedicates 32 slots to each stream. Each slot holds 64 bytes of data.

(158) Butterfly network **2717/2727** consists of a 7 stage butterfly network. Butterfly network **2717/2727** receives 128 bytes of input and generates 64 bytes of output. The first stage of the butterfly is actually a half-stage. It collects bytes from both slots that match a non-aligned fetch and merges them into a single, rotated 64-byte array. The remaining 6 stages form a standard butterfly network. Butterfly network **2717/2727** performs the following operations: rotates the next element down to byte lane **0**; promotes data types by one power of 2, if requested; swaps real and imaginary components of complex numbers, if requested; converts big endian to little endian if central processing unit core **110** is presently in big endian mode. The user specifies element size, type promotion and real/imaginary swap as part of the stream's parameters.

(159) Streaming engine **2700** attempts to fetch and format data ahead of central processing unit core **110**'s demand for it, so that it can maintain full throughput. Holding registers **2718/2728** provide a small amount of buffering so that the process remains fully pipelined. Holding registers **2718/2728** are not directly architecturally visible, except for the fact that streaming engine **2700** provides full throughput.

(160) The two streams **2710/2720** share a pair of independent L2 interfaces **2730**: L2 Interface A (IFA) **2733** and L2 Interface B (IFB) **2734**. Each L2 interface provides 512 bits/cycle throughput direct to the L2 controller for an aggregate bandwidth of 1024 bits/cycle. The L2 interfaces use the credit-based multicore bus architecture (MBA) protocol. The L2 controller assigns each interface its own pool of command credits. The pool should have sufficient credits so that each interface can send sufficient requests to achieve full read-return bandwidth when reading L2 RAM, L2 cache and multicore shared memory controller (MSMC) memory (described below).

(161) To maximize performance, both streams can use both L2 interfaces, allowing a single stream to send a peak command rate of 2 requests/cycle. Each interface prefers one stream over the other, but this preference changes dynamically from request to request. IFA **2733** and IFB **2734** always prefer opposite streams, when IFA **2733** prefers Stream **0**, IFB **2734** prefers Stream **1** and vice versa.

(162) Arbiter **2731/2732** ahead of each interface **2733/2734** applies the following basic protocol on every cycle it has credits available. Arbiter **2731/2732** checks if the preferred stream has a command ready to send. If so, arbiter **2731/2732** chooses that command. Arbiter **2731/2732** next checks if an alternate stream has at least two requests ready to send, or one command and no credits. If so, arbiter **2731/2732** pulls a command from the alternate stream. If either interface issues a command, the notion of preferred and alternate streams swap for the next request. Using this simple algorithm, the two interfaces dispatch requests as quickly as possible while retaining fairness between the two streams. The first rule ensures that each stream can send a request on every cycle that has available credits. The second rule provides a mechanism for one stream to borrow the other's interface when the second interface is idle. The third rule spreads the bandwidth demand for each stream across both interfaces, ensuring neither interface becomes a bottleneck by itself.

(163) Coarse Grain Rotator **2735/2736** enables streaming engine **2700** to support a transposed matrix addressing mode. In this mode, streaming engine **2700** interchanges the two innermost dimensions of its multidimensional loop. This accesses an array column-wise rather than row-wise. Rotator **2735/2736** is not architecturally visible, except as enabling this transposed access mode.

(164) The stream definition template provides the full structure of a stream that contains data. The iteration counts and dimensions provide most of the structure, while the various flags provide the rest of the details. For all data-containing streams, the streaming engine defines a single stream template. All stream types it supports fit this template. The streaming engine defines a six-level loop nest for addressing elements within the stream. Most of the fields in the stream template map directly to the parameters in that algorithm. FIG. **28** illustrates stream template register **2800**. The numbers above the fields are bit numbers within a 256-bit vector. Table 6 shows the stream field definitions of a stream template.

(165) TABLE-US-00006

TABLE 6	FIG. 28	Field	Reference	Size	Name	Number	Description	Bits	
ICNT0	2801	Iteration count for loop 0	16	ICNT1	2802	Iteration count for loop 1	16	ICNT2	2803
Iteration count for loop 2	16	ICNT3	2804	Iteration count for loop 3	16	ICNT4	2805	Iteration count for loop 4	16
INCT5	2806	Iteration count for loop 5	16	DIM1	2822	Signed dimension for loop 1	16	DIM2	2823
Signed dimension for loop 2	16	DIM3	2824	Signed dimension for loop 3	16	DIM4	2825	Signed dimension for loop 4	32
DIM5	2826	Signed dimension for loop 5	32	FLAGS	2811	Stream modifier flags	48		

(166) Loop 0 is the innermost loop and loop 5 is the outermost loop. In the current example DIM0 is always equal to is ELEM_BYTES defining physically contiguous data. Thus the stream template register **2800** does not define DIM0. Streaming engine **2700** interprets all iteration counts as unsigned integers and all dimensions as unscaled signed integers. The template above fully specifies the type of elements, length and dimensions of the stream. The stream instructions separately specify a start address. This would typically be by specification of a scalar register in scalar register file **211** which stores this start address. This allows a program to open multiple streams using the same template.

(167) FIG. **29** illustrates sub-field definitions of the flags field **2811**. As shown in FIG. **29** the flags field **2811** is 6 bytes or 48 bits. FIG. **29** shows bit numbers of the fields. Table 7 shows the definition of these fields.

(168) TABLE-US-00007

TABLE 7	FIG. 29	Reference	Field	Name	Number	Description	Bits
ELTYPE	2901	Type of data element	4	TRANPOSE	2902	Two dimensional transpose mode	3
PROMOTE	2903	Promotion mode	3	VCLLEN	2904	Stream vector length	3
ELDUP	2905	Element					

duplication 3 GRDUP 2906 Group duplication 1 DECIM 2907 Element decimation 2 THROTTLE 2908 Fetch ahead throttle mode 2 DIMFMT 2909 Stream dimensions format 3 DIR 2910 Stream direction 1 0 forward direction 1 reverse direction CBK0 2911 First circular block size number 4 CBK1 2912 Second circular block size number 4 AM0 2913 Addressing mode for loop 0 2 AM1 2914 Addressing mode for loop 1 2 AM2 2915 Addressing mode for loop 2 2 AM3 2916 Addressing mode for loop 3 2 AM4 2917 Addressing mode for loop 4 2 AM5 2918 Addressing mode for loop 5 2

(169) The Element Type (ELTYPE) field **2901** defines the data type of the elements in the stream. The coding of the four bits of the ELTYPE field **2901** is defined as shown in Table 8.

(170) TABLE-US-00008 TABLE 8 Sub-element Total Element ELTYPE Real/Complex Size Bits
 Size Bits 0000 real 8 8 0001 real 16 16 0010 real 32 32 0011 real 64 64 0100 reserved 0101 reserved 0110 reserved 0111 reserved 1000 complex no swap 8 16 1001 complex no swap 16 32 1010 complex no swap 32 64 1011 complex no swap 64 128 1100 complex swapped 8 16 1101 complex swapped 16 32 1110 complex swapped 32 64 1111 complex swapped 64 128

(171) Real/Complex Type determines whether the streaming engine treats each element as a real number or two parts (real/imaginary or magnitude/angle) of a complex number. This field also specifies whether to swap the two parts of complex numbers. Complex types have a total element size that is twice their sub-element size. Otherwise, the sub-element size equals total element size.

(172) Sub-Element Size determines the type for purposes of type promotion and vector lane width. For example, 16-bit sub-elements get promoted to 32-bit sub-elements when a stream requests type promotion. The vector lane width matters when central processing unit core **110** operates in big endian mode, as it always lays out vectors in little endian order.

(173) Total Element Size determines the minimal granularity of the stream. In the stream addressing model, it determines the number of bytes the stream fetches for each iteration of the innermost loop. Streams always read whole elements, either in increasing or decreasing order. Therefore, the innermost dimension of a stream spans $ICNT0 \times \text{total-element-size}$ bytes.

(174) The TRANSPOSE field **2902** determines whether the streaming engine accesses the stream in a transposed order. The transposed order exchanges the inner two addressing levels. The TRANSPOSE field **2902** also indicated the granularity it transposes the stream. The coding of the three bits of the TRANSPOSE field **2902** is defined as shown in Table 9 for normal 2D operations.

(175) TABLE-US-00009 TABLE 9 Transpose Meaning 000 Transpose disabled 001 Transpose on 8-bit boundaries 010 Transpose on 16-bit boundaries 011 Transpose on 32-bit boundaries 100 Transpose on 64-bit boundaries 101 Transpose on 128-bit boundaries 110 Transpose on 256-bit boundaries 111 Reserved

(176) Streaming engine **2700** may transpose data elements at a different granularity than the element size. This allows programs to fetch multiple columns of elements from each row. The transpose granularity must be no smaller than the element size. The TRANSPOSE field **2902** interacts with the DIMFMT field **2909** in a manner further described below.

(177) The PROMOTE field **2903** controls whether the streaming engine promotes sub-elements in the stream and the type of promotion. When enabled, streaming engine **2700** promotes types by a powers-of-2 sizes. The coding of the three bits of the PROMOTE field **2903** is defined as shown in Table 10.

(178) TABLE-US-00010 TABLE 10 Promotion Promotion Resulting Sub-element Size
 PROMOTE Factor Type 8-bit 16-bit 32-bit 64-bit 000 $1 \times$ N/A 8-bit 16-bit 32-bit 64-bit 001 $2 \times$ zero 16-bit 32-bit 64-bit Invalid extend 010 $4 \times$ zero 32-bit 64-bit Invalid Invalid extend 011 $8 \times$ zero 64-bit Invalid Invalid Invalid extend 100 reserved 101 $2 \times$ sign 16-bit 32-bit 64-bit Invalid extend 110 $4 \times$ sign 32-bit 64-bit Invalid Invalid extend 111 $8 \times$ sign 64-bit Invalid Invalid Invalid extend

(179) When PROMOTE is 000, corresponding to a $1 \times$ promotion, each sub-element is unchanged and occupies a vector lane equal in width to the size specified by ELTYPE. When PROMOTE is 001, corresponding to a $2 \times$ promotion and zero extend, each sub-element is treated as an unsigned

integer and zero extended to a vector lane twice the width specified by ELTYPE. A 2× promotion is invalid for an initial sub-element size of 64 bits. When PROMOTE is 010, corresponding to a 4× promotion and zero extend, each sub-element is treated as an unsigned integer and zero extended to a vector lane four times the width specified by ELTYPE. A 4× promotion is invalid for an initial sub-element size of 32 or 64 bits. When PROMOTE is 011, corresponding to an 8× promotion and zero extend, each sub-element is treated as an unsigned integer and zero extended to a vector lane eight times the width specified by ELTYPE. An 8× promotion is invalid for an initial sub-element size of 16, 32 or 64 bits. When PROMOTE is 101, corresponding to a 2× promotion and sign extend, each sub-element is treated as a signed integer and sign extended to a vector lane twice the width specified by ELTYPE. A 2× promotion is invalid for an initial sub-element size of 64 bits. When PROMOTE is 110, corresponding to a 4× promotion and sign extend, each sub-element is treated as a signed integer and sign extended to a vector lane four times the width specified by ELTYPE. A 4× promotion is invalid for an initial sub-element size of 32 or 64 bits. When PROMOTE is 111, corresponding to an 8× promotion and zero extend, each sub-element is treated as a signed integer and sign extended to a vector lane eight times the width specified by ELTYPE. An 8× promotion is invalid for an initial sub-element size of 16, 32 or 64 bits.

(180) The VECLLEN field **2904** defines the stream vector length for the stream in bytes. Streaming engine **2700** breaks the stream into groups of elements that are VECLLEN bytes long. The coding of the three bits of the VECLLEN field **2904** is defined as shown in Table 11.

(181) TABLE-US-00011 TABLE 11 VECLLEN Stream Vector Length 000 1 byte 001 2 bytes 010 4 bytes 011 8 bytes 100 16 bytes 101 32 bytes 110 64 bytes 111 Reserved

(182) VECLLEN must be greater than or equal to the product of the element size in bytes and the duplication factor. Streaming engine **2700** presents the stream to central processing unit core **110** as either a sequence of pairs of single vectors or a sequence of double vectors. When VECLLEN is shorter the native vector width of central processing unit core **110**, streaming engine **2700** pads the extra lanes in the vector provided to central processing unit core **110**. The GRDUP field **2906** determines the type of padding. The VECLLEN field **2904** interacts with ELDUP field **2905** and GRDUP field **2906** in a manner detailed below.

(183) The ELDUP field **2905** specifies a number of times to duplicate each element. The element size multiplied with the element duplication amount must not exceed the 64 bytes. The coding of the three bits of the ELDUP field **2905** is defined as shown in Table 12.

(184) TABLE-US-00012 TABLE 12 ELDUP Duplication Factor 000 No Duplication 001 2 times 010 4 times 011 8 times 100 16 times 101 32 times 110 64 times 111 Reserved

The ELDUP field **2905** interacts with VECLLEN field **2904** and GRDUP field **2906** in a manner detailed below.

(185) The GRDUP bit **2906** determines whether group duplication is enabled. If GRDUP bit **2906** is 0, then group duplication is disabled. If the GRDUP bit **2906** is 1, then group duplication is enabled. When enabled by GRDUP bit **2906**, streaming engine **2700** duplicates a group of elements to fill the vector width. VECLLEN field **2904** defines the length of the group to replicate. When VECLLEN field **2904** is less than the vector length of central processing unit core **110** and GRDUP bit **2906** enables group duplication, streaming engine **2700** fills the extra lanes (see FIGS. **20** and **21**) with additional copies of the stream vector. Because stream vector lengths and vector length of central processing unit core **110** are always powers of two, group duplication always produces a power of two of the number of duplicate copies. GRDUP field **2906** specifies how stream engine **2700** pads stream vectors out to the vector length of central processing unit core **110**. When GRDUP bit **2906** is 0, streaming engine **2700** fills the extra lanes with zero and marks these extra vector lanes invalid. When GRDUP bit **2906** is 1, streaming engine **2700** fills extra lanes with copies of the group of elements in each stream vector. Setting GRDUP bit **2906** to **1** has no effect when VECLLEN is set to the native vector width of central processing unit core **110**.

(186) The DECIM field **2907** controls data element decimation of the corresponding stream.

Streaming engine **2700** deletes data elements from the stream upon storage in head registers **2718/2728** for presentation to the requesting functional unit. Decimation always removes whole data elements, not sub-elements. The DECIM field **2907** is defined as listed in Table 13.

(187) TABLE-US-00013 TABLE 13 DECIM Decimation Factor 00 No Decimation 01 2 times 10 4 times 11 Reserved

(188) If DECIM field **2907** equals 00, then no decimation occurs. The data elements are passed to the corresponding head registers **2718/2728** without change. If DECIM field **2907** equals 01, then 2:1 decimation occurs. Streaming engine **2700** removes odd number elements from the data stream upon storage in the head registers **2718/2728**. Limitations in the formatting network require 2:1 decimation to be employed with data promotion by at least $2\times$ (PROMOTE cannot be 000), ICNT0 must be multiple of 2 and the total vector length (VECLen) must be large enough to hold a single promoted, duplicated element. For transposed streams (TRANSPPOSE $\neq$ 0), the transpose granule must be at least twice the element size in bytes before promotion. If DECIM field **2907** equals 10, then 4:1 decimation occurs. Streaming engine **2700** retains every fourth data element removing three elements from the data stream upon storage in the head registers **2718/2728**. Limitations in the formatting network require 4:1 decimation to be employed with data promotion by at least $4\times$ (PROMOTE cannot be 000, 001 or 101), ICNT0 must be multiple of 4 and the total vector length (VECLen) must be large enough to hold a single promoted, duplicated element. For transposed streams (TRANSPPOSE $\neq$ 0), decimation always removes columns, and never removes rows. Thus the transpose granule must be: at least twice the element size in bytes before promotion for 2:1 decimation ($GRANULE \geq 2 \times ELEM_BYTES$); and at least four times the element size in bytes before promotion for 4:1 decimation ($GRANULE \geq 4 \times ELEM_BYTES$).

(189) The THROTTLE field **2908** controls how aggressively the streaming engine fetches ahead of central processing unit core **110**. The coding of the two bits of this field is defined as shown in Table 14.

(190) TABLE-US-00014 TABLE 14 THROTTLE Description 00 Minimum throttling, maximum fetch ahead 01 Less throttling, more fetch ahead 10 More throttling, less fetch ahead 11 Maximum throttling, minimum fetch ahead

(191) THROTTLE does not change the meaning of the stream, and serves only as a hint. The streaming engine may ignore this field. Programs should not rely on the specific throttle behavior for program correctness, because the architecture does not specify the precise throttle behavior. THROTTLE allows programmers to provide hints to the hardware about the program's own behavior. By default, the streaming engine attempts to get as far ahead of central processing unit core **110** as it can to hide as much latency as possible, while providing full stream throughput to central processing unit core **110**. While several key applications need this level of throughput, it can lead to bad system level behavior for others. For example, the streaming engine discards all fetched data across context switches. Therefore, aggressive fetch-ahead can lead to wasted bandwidth in a system with large numbers of context switches. Aggressive fetch-ahead only makes sense in those systems if central processing unit core **110** consumes data very quickly.

(192) The DIMFMT field **2909** enables redefinition of the loop count fields ICNT0 **2801**, ICNT1 **2802**, ICNT2 **2803**, ICNT3 **2804**, ICNT4 **2805** and ICNT5 **2806**, the loop dimension fields DIM **2855**, DIM2 **2823**, DIM3 **2824**, DIM4 **2825** and DIM5 **2826** and the addressing mode fields AM0 **2913**, AM1 **2914**, AM2 **2915**, AM3 **2916**, AM4 **2917** and AM5 **2918** (part of FLAGS field **2811**) of the stream template register **2800**. This permits some loop dimension fields and loop counts to include more bits at the expense of fewer loops. Table 15 lists the size of the loop dimension fields for various values of the DIMFMT field **2909**.

(193) TABLE-US-00015 TABLE 15 Number DIMFMT of Loops DIM5 DIM4 DIM3 DIM2 DIM1

000	3 unused	32 bits unused	32 bits unused	001	4 unused	32 bits unused	16 bits	16 bits	010	4 unused	32 bits	16 bits	16 bits unused
011	5 unused	32 bits	32 bits	32 bits	16 bits	100	reserved	101	reserved	110	6	16 bits	16 bits
111	6	32 bits	32 bits	16 bits	16 bits	32 bits	32 bits	111	6	32 bits	32 bits	16 bits	16 bits

(194) Note that DIM0 always equals ELEM_BYTES. Table 16 lists the size of the loop count fields for various values of the DIMFMT field **2909**.

(195) TABLE-US-00016 TABLE 16 Number DIMFMT of Loops ICNT5 ICNT4 ICNT3 ICNT2 ICNT1 ICNT0 000 3 unused 32 bits unused 32 bits unused 32 bits 001 4 unused 32 bits unused 32 bits 16 bits 16 bits 010 4 unused 32 bits 16 bits 16 bits unused 32 bits 011 5 unused 32 bits 16 bits 16 bits 16 bits 100 reserved 101 reserved 110 6 16 bits 16 bits 16 bits 16 bits 16 bits 111 6 16 bits 16 bits 16 bits 16 bits 16 bits 16 bits

DIMFMT field **2909** effectively defines the loop dimension and loop count bits of stream template register **2800**. FIG. **28** illustrates the default case when DIMFMT is 111.

(196) FIGS. **30** to **34** illustrate the definition of bits of the stream template register for other values of DIMFMT. Note the location and meaning of the FLAGS field (**2811**, **3011**, **3111**, **3211**, **3311** and **3411**) are the same for all values of DIMFMT

(197) FIG. **30** illustrates the definition of bits of the stream template register **3000** for a DIMFMT value of 000. For a DIMFMT value of 000, there are three loops: loop0, loop2 and loop4. For loop0 ICNT0 field **3001** includes bits **0** to **31** and DIM0 field equals ELEM_BYTES. For loop2 ICNT2 field **3002** includes bits **32** to **63** and DIM2 field **3021** includes bits **160** to **191**. For loop4 ICNT4 field **3003** includes bits **64** to **95** and DIM4 field **3022** includes bits **192** to **223**.

(198) FIG. **31** illustrates the definition of bits of the stream template register **3100** for a DIMFMT value of 001. For a DIMFMT value of 001, there are four loops: loop0, loop1, loop2 and loop4. For loop0 ICNT0 field **3101** includes bits **0** to **16** and DIM0 field equals ELEM_BYTES. For loop1 ICNT1 field **3002** includes bits **16** to **31** and DIM1 field **3123** includes bits **224** to **255**. For loop2 ICNT2 field **3103** includes bits **32** to **63** and DIM2 field **3121** includes bits **160** to **191**. For loop4 ICNT4 field **3104** includes bits **64** to **95** and DIM4 field **3122** includes bits **192** to **223**.

(199) FIG. **32** illustrates the definition of bits of the stream template register **3200** for a DIMFMT value of 010. For a DIMFMT value of 010, there are four loops: loop0, loop2, loop3 and loop4. For loop0 ICNT0 field **3201** includes bits **0** to **32** and DIM0 field equals ELEM_BYTES. For loop2 ICNT2 field **3202** includes bits **32** to **47** and DIM2 field **3221** includes bits **160** to **191**. For loop3 ICNT3 field **3203** includes bits **48** to **63** and DIM3 field **3223** includes bits **224** to **255**. For loop4 ICNT4 field **3204** includes bits **64** to **95** and DIM4 field **3222** includes bits **192** to **223**.

(200) FIG. **33** illustrates the definition of bits of the stream template register **3300** for a DIMFMT value of 011. For a DIMFMT value of 011, there are five loops: loop0, loop1, loop2, loop3 and loop4. For loop0 ICNT0 field **3401** includes bits **0** to **15** and DIM0 field equals ELEM_BYTES. For loop1 ICNT1 field **3402** includes bits **16** to **31** and DIM1 field **3421** includes bits **144** to **159**. For loop2 ICNT2 field **3403** includes bits **32** to **47** and DIM2 field **3221** includes bits **160** to **191**. For loop3 ICNT3 field **3204** includes bits **48** to **63** and DIM3 field **3424** includes bits **224** to **255**. For loop4 ICNT4 field **3405** includes bits **64** to **95** and DIM4 field **3423** includes bits **192** to **223**.

(201) FIG. **34** illustrates the definition of bits of the stream template register **3400** for a DIMFMT value of 101. For a DIMFMT value of 110, there are six loops: loop0, loop1, loop2, loop3, loop4 and loop5. For loop0 ICNT0 field **3501** includes bits **0** to **15** and DIM0 field equals ELEM_BYTES. For loop1 ICNT1 field **3502** includes bits **16** to **31** and DIM1 field **3521** includes bits **144** to **159**. For loop2 ICNT2 field **3503** includes bits **32** to **47** and DIM2 field **3522** includes bits **160** to **191**. For loop3 ICNT3 field **3504** includes bits **48** to **63** and DIM3 field **3525** includes bits **224** to **255**. For loop4 ICNT4 field **3405** includes bits **64** to **79** and DIM4 field **3523** includes bits **192** to **207**. For loop5 ICNT5 field **3506** includes bits **80** to **95** and DIM5 field **3524** includes bits **208** to **223**.

(202) The DIR bit **2910** determines the direction of fetch of the inner loop (Loop0). If the DIR bit **2910** is 0 then Loop0 fetches are in the forward direction toward increasing addresses. If the DIR bit **2910** is 1 then Loop0 fetches are in the backward direction toward decreasing addresses. The fetch direction of other loops is determined by the sign of the corresponding loop dimension DIM1, DIM2, DIM3, DIM4 and DIM5 which are signed integers.

(203) The CBK0 field **2911** and the CBK1 field **2912** control the circular block size upon selection of circular addressing. The manner of determining the circular block size will be more fully described below.

(204) The AM0 field **2913**, AM1 field **2914**, AM2 field **2915**, AM3 field **2916**, AM4 field **2917** and AM5 field **2918** control the addressing mode of a corresponding loop. This permits the addressing mode to be independently specified for each loop. Each of AM0 field **2913**, AM1 field **2914**, AM2 field **2915**, AM3 field **2916**, AM4 field **2917** and AM5 field **2918** are three bits and are decoded as listed in Table 17.

(205) TABLE-US-00017

TABLE 17 AMx field Meaning	
00	Linear addressing
01	Circular addressing
10	Circular addressing block size set by CBK0
11	Circular addressing block size set by CBK0 + CBK1 + 1
11	reserved

(206) In linear addressing the address advances according to the address arithmetic whether forward or reverse. In circular addressing the address remains within a defined address block. Upon reaching the end of the circular address block the address wraps around to other limit of the block. Circular addressing blocks are typically limited to $2^{\text{sup.}N}$ addresses where N is an integer. Circular address arithmetic may operate by cutting the carry chain between bits and not allowing a selected number of most significant bits to change. Thus arithmetic beyond the end of the circular block changes only the least significant bits.

(207) The block size is set as listed in Table 18.

(208) TABLE-US-00018

TABLE 18 Encoded Block Size CBK0 or Block Size CBK0 + CBK1 + 1 (bytes)	
0	512
1	1K
2	2K
3	4K
4	8K
5	16K
6	32K
7	64K
8	128K
9	256K
10	512K
11	1M
12	2M
13	4M
14	8M
15	16M
16	32M
17	64M
18	128M
19	256M
20	512M
21	1G
22	2G
23	4G
24	8G
25	16G
26	32G
27	64G
28	Reserved
29	Reserved
30	Reserved
31	Reserved

(209) In the preferred embodiment the circular block size is set by the number encoded by CBK0 (first circular address mode 01) or the number encoded by CBK0+CBK1+1 (second circular address mode 10). For example the first circular address mode, the circular address block size can be from 512 bytes to 16 M bytes. For the second circular address mode, the circular address block size can be from 1 K bytes to 64 G bytes. Thus the encoded block size is $2^{\text{sup.}(B+9)}$ bytes, where B is the encoded block number which is CBK0 for the first block size (AMx of 01) and CBK0+CBK1+1 for the second block size (AMx of 10).

(210) FIG. 35 illustrates loop count selection circuit **3500** which is an exemplary embodiment selecting data from the stream template register for the various loop dimensions. As illustrated in FIGS. 28 and 30 to 34 the stream template register bits defining the loop counts vary dependent upon the DIMFMT field. FIG. 35 illustrates bits **0** to **95** of the stream template register. These bits are divided into 6 portions including: portion **3501**, bits **0** to **15**; portion **3502**, bits **16** to **31**; portion **3503**, bits **32** to **47**; portion **3504**, bits **48** to **63**; portion **3505**, bits **64** to **79**; and portion **3506**, bits **80** to **95**.

(211) Concatenator **3511** forms a single 32-bit data word from portions **3501** and **3502**. Multiplexer **3512** selects either portion **3501** or the output of concatenator **3511** for the INCT0 output. Multiplexer **3513** selects either a null input or portion **3502** for the INCT1 output.

(212) Concatenator **3521** forms a single 32-bit data word from portions **3503** and **3504**. Multiplexer **3522** selects either portion **3503** or the output of concatenator **3521** for the INCT2 output. Multiplexer **3523** selects either a null input or portion **3504** for the INCT3 output.

(213) Concatenator **3531** forms a single 32-bit data word from portions **3505** and **3506**. Multiplexer **3532** selects either portion **3505** or the output of concatenator **3531** for the INCT4 output. Multiplexer **3533** selects either a null input or portion **3506** for the INCT5 output.

(214) DIMFMT ICNT decoder **3508** receives the DIMFT bits from the stream template register and generates outputs controlling the selections of multiplexers **3512**, **3513**, **3522**, **3523**, **3532** and **3533**. Table 19 lists the control of these multiplexers for the various codings of the DIMFMT field.

(215) TABLE-US-00019

TABLE 19	ICNT5	ICNT4	ICNT3	ICNT2	ICNT1	ICNT0	MUX	MUX
MUX MUX MUX MUX DIMFMT	3533	3532	3523	3522	3513	3512	000	Null
3531	Null	3531	Null	3521	3502	3501	010	Null
3531	3504	3503	Null	3511	011	Null	3531	
3504	3503	3502	3501	100	reserved	101	reserved	110
3506	3505	3504	3503	3502	3501	111	3506	
3505	3504	3503	3502	3501				

(216) FIG. 36 illustrates loop dimension selection circuit 3600 which is an exemplary embodiment selecting data from the stream template register for the various loop dimensions. Note that DIM0, the loop dimension of loop0, is always ELEN_BYTES. As illustrated in FIGS. 28 and 30 to 34 the stream template register bits defining the loop dimension vary dependent upon the DIMFMT field. FIG. 36 illustrates bits 144 to 255 of the stream template register. These bits are divided into 6 portions including: portion 3601, bits 144 to 159; portion 3602, bits 160 to 175; portion 3603, bits 176 to 191; portion 3604, bits 192 to 207; portion 3605, bits 208 to 223; and portion 3606, bits 224 to 255.

(217) Concatenator 3611 forms a single 32-bit data word from portions 3602 and 3603. Multiplexer 3612 selects either portion 3601, a null input or portion 3606 a DIM1 output. Multiplexer 3613 selects either portion 3602 of the output of concatenator 3611 for the DIM2 output. Multiplexer 3614 selects either portion 3604, a null input or portion 3606 for the DIM3 output.

(218) Concatenator 3621 forms a single 32-bit data word from portions 3604 and 3605. Multiplexer 3622 selects either portion 3604 or the output of concatenator 3621 for the DIM4 output. Multiplexer 3623 selects either a null input, portion 3605 or portion 3606 for the DIM5 output.

(219) DIMFMT DIM decoder 3607 receives the DIMFT bits from the stream template register and generates outputs controlling the selections of multiplexers 3612, 3613, 3614, 3622 and 3623. Table 20 lists the control of these multiplexers for the various codings of the DIMFMT field.

(220) TABLE-US-00020

TABLE 20	DIM5	DIM4	DIM3	DIM2	DIM1	MUX	MUX	MUX	MUX
MUX DIMFMT	3623	3622	3614	3613	3612	000	Null	3621	Null
3611	3606	010	Null	3621	3606	3611	Null	011	Null
3621	3606	3611	3601	100	reserved	101	reserved	110	3605
3604	3606	3611	3601	111	3606	3621	3603	3602	3601

(221) FIG. 37 illustrates an example of adder control word circuit 3700 which generates an adder control word for loop address generators to be described below. The addressing mode fields AM0 2913, AM1 2914, AM2 2915, AM3 2916, AM4 2917 and AM5 2918 each control the addressing mode of a corresponding loop of the stream engine address. The address control word circuit 3700 is provided for each supported loop of the streaming engine address. Adder 3701 forms the sum of CBK0, CBK1 and 1. The fields CBK0 and CBK1 are each 4-bit fields and are part of the EFLAGS field of the corresponding stream template register. The fields CBK0 and CBK1 are supplied to the operand inputs of adder 3701. The quantity +1 is supplied to the carry-input of the least significant bit of adder 3701. This structure enables the three term addition without special adder hardware. The sum output of adder 3701 supplies one input of multiplexer 3702. A second input of multiplexer 3702 is a null. A third input of multiplexer 3702 is CBK0. Multiplexer 3702 is controlled by the AMx field for the corresponding loop. If AMx is 000, then multiplexer 3702 selects the null input. If AMx is 001, then multiplexer 3702 selects the CBK0 input. If AMx is 010, then multiplexer 3702 selects the sum CBK0+CBK1+1 output from adder 3701. The output of multiplexer 3702 is used as an index into adder control word look up table 3703. The adder control words accessed from adder control word look up table 3703 are used to control a corresponding loop adder (see FIG. 38) in manner similar to the SIMD control described above in conjunction with FIG. 18 and Table 3. The corresponding loop adder includes carry break circuits as illustrated in FIG. 18 following bits corresponding to the supported block sizes listed in Table 18. Adder control word look up table 3703 includes control words such as listed in Table 3 for the various block sizes. The selected number CBK0 or CBK0+CBK1+1 indexes the appropriate adder control word. If multiplexer 3702 selects the null input, corresponding to linear addressing, the corresponding adder control word is all 1's permitting adder carries between all address bits.

(222) FIG. 38 illustrates a partial schematic view of a streaming engine 2700 address generator 3800. Address generator 3800 forms an address for fetching a next element in the defined stream of the corresponding streaming engine. Start address register 3801 stores a start address of the data stream. As previously described, start address register 3801 is preferably a scalar register in global scalar register file 211 designated by the STROPEN instruction that opened the corresponding stream. As known in the art, this start address may be copied from the specified scalar register and stored locally at the corresponding address generator 2711 OR 2721. A first loop of the stream employs Loop0 count register 3811, adder 3812, multiplier 3813 and comparator 3814. Loop0 count register 3811 stores the working copy of the iteration count of the first loop (Loop0). For each iteration of Loop0 adder 3812, as triggered by the Next Address signal, adds 1 to the loop count, which is stored back in Loop0 count register 3811. Multiplier 3813 multiplies the current loop count and the quantity ELEM_BYTES. ELEM_BYTES is the size of each data element in the loop in bytes. Loop0 traverses data elements physically contiguous in memory of the iteration step size is ELEM_BYTES.

(223) Comparator 3814 compares the count stored in Loop0 count register 3811 (after incrementing by adder 3812) with the value of ICNT0 2801 from the corresponding stream template register 2800. As described above the loop0 count may be at portion 3001, portion 3101, portion 3201, portion 3301 or portion 3401 depending upon the state of the DIMFMT field 2909. When the output of adder 3812 equals the value of ICNT0 2801 of the stream template register 2800, an iteration of Loop0 is complete. Comparator 3814 generates an active Loop0 End signal. Loop0 count register 3811 is reset to 0 and an iteration of the next higher loop, in this case Loop1, is triggered.

(224) Circuits for the higher loops (Loop1, Loop2, Loop3, Loop4, Loop5) are similar to that illustrated in FIG. 38. Each loop includes a corresponding working loop count register, adder, multiplier and comparator. The adder of each loop is triggered by the loop end signal of the prior loop. The second input to each multiplier is the corresponding dimension DIM1, DIM2, DIM3, DIM4 and DIM5 of the corresponding stream template. The comparator of each loop compares the working loop register count with the corresponding iteration value ICNT1, ICTN2, ICTN3, ICTN4 and ICTN5 of the corresponding stream template register. A loop end signal generates an iteration of the next higher loop. A loop end signal from loop5 ends the stream.

(225) FIG. 38 illustrates adder 3812 receiving an adder control word. As described above this adder control word is all 1's for linear addressing and has a 0 at the appropriate location for circular addressing. The location of the 0 in the adder control word corresponding to the circular block size of the circular addressing mode.

(226) The central processing unit core 110 exposes the streaming engine to programs through a small number of instructions and specialized registers. A STROPEN instruction opens a stream. The STROPEN command specifies a stream number indicating opening stream 0 or stream 1. The STROPEN specifies a stream template register which stores the stream template as described above. The arguments of the STROPEN instruction are listed in Table 21.

(227) TABLE-US-00021 TABLE 21 Argument Description Stream Start Scaler register storing stream start Address Register address Steam Number Stream 0 or Stream 1 Stream Template Vector register storing stream Register template data

(228) The stream start address register is preferably a scalar register in general scalar register file 211. The STROPEN instruction specifies stream 0 or stream 1 by its opcode. The stream template register is preferably a vector register in general vector register file 221. If the specified stream is active the STROPEN instruction closes the prior stream and replaces the stream with the specified stream.

(229) A STRCLOSE instruction closes a stream. The STRCLOSE command specifies the stream number of the stream to be closed.

(230) A STRSAVE instruction captures sufficient state information of a specified stream to restart

that stream in the future. A STRRSTR instruction restores a previously saved stream. A STRSAVE instruction does not save any of the data of the stream. A STRSAVE instruction saves only metadata. The stream re-fetches data in response to a STRRSTR instruction.

(231) Streaming engine is in one of three states: Inactive; Active; or Frozen. When inactive the streaming engine does nothing. Any attempt to fetch data from an inactive streaming engine is an error. Until the program opens a stream, the streaming engine is inactive. After the program consumes all the elements in the stream or the program closes the stream, the streaming engine also becomes inactive. Programs which use streams explicitly activate and inactivate the streaming engine. The operating environment manages streams across context-switch boundaries via the streaming engine's implicit freeze behavior, coupled with its own explicit save and restore actions.

(232) Active streaming engines have a stream associated with them. Programs can fetch new stream elements from active streaming engines. Streaming engines remain active until one of the following. When the stream fetches the last element from the stream, it becomes inactive. When program explicitly closes the stream, it becomes inactive. When central processing unit core **110** responds to an interrupt or exception, the streaming engine freezes. Frozen streaming engines capture all the state necessary to resume the stream where it was when the streaming engine froze. The streaming engines freeze in response to interrupts and exceptions. This combines with special instructions to save and restore the frozen stream context, so that operating environments can cleanly switch contexts. Frozen streams reactivate when central processing unit core **110** returns to the interrupted context.

(233) FIG. **39** is a partial schematic diagram **3900** illustrating the stream input operand coding described above. FIG. **39** illustrates decoding src1 field **1305** of one instruction of a corresponding src1 input of functional unit **3920**. These same circuits are duplicated for src2/cst field **1304** and the src2 input of functional unit **3920**. In addition, these circuits are duplicated for each instruction within an execute packet that can be dispatched simultaneously.

(234) Instruction decoder **113** receives bits **13** to **17** comprising src1 field **1305** of an instruction. The opcode field opcode field (bits **4** to **12** for all instructions and additionally bits **28** to **31** for unconditional instructions) unambiguously specifies a corresponding functional unit **3920**. In this embodiment functional unit **3920** could be L2 unit **241**, S2 unit **242**, M2 unit **243**, N2 unit **244** or C unit **245**. The relevant part of instruction decoder **113** illustrated in FIG. **39** decodes src1 bit field **1305**. Sub-decoder **3911** determines whether src1 bit field **1305** is in the range from 00000 to 01111. If this is the case, sub-decoder **3911** supplies a corresponding register number to global vector register file **231**. In this example this register field is the four least significant bits of src1 bit field **1305**. Global vector register file **231** recalls data stored in the register corresponding to this register number and supplies this data to the src1 input of functional unit **3920**. This decoding is generally known in the art.

(235) Sub-decoder **3912** determines whether src1 bit field **1305** is in the range from 10000 to 10111. If this is the case, sub-decoder **3912** supplies a corresponding register number to the corresponding local vector register file. If the instruction is directed to L2 unit **241** or S2 unit **242**, the corresponding local vector register file is local vector register field **232**. If the instruction is directed to M2 unit **243**, N2 unit **244** or C unit **245**, the corresponding local vector register file is local vector register field **233**. In this example this register field is the three least significant bits of src1 bit field **1305**. The corresponding local vector register file **232/233** recalls data stored in the register corresponding to this register number and supplies this data to the src1 input of functional unit **3920**. This decoding is generally known in the art.

(236) Sub-decoder **3913** determines whether src1 bit field **1305** is 11100. If this is the case, sub-decoder **3913** supplies a stream 0 read signal to streaming engine **2700**. Streaming engine **2700** then supplies stream 0 data stored in holding register **2718** to the src1 input of functional unit **3920**.

(237) Sub-decoder **3914** determines whether src1 bit field **1305** is 11101. If this is the case, sub-decoder **3914** supplies a stream 0 read signal to streaming engine **2700**. Streaming engine **2700**

then supplies stream 0 data stored in holding register **2718** to the src1 input of functional unit **3920**. Sub-decoder **3914** also supplies an advance signal to stream 0. As previously described, streaming engine **2700** advances to store the next sequential data elements of stream 0 in holding register **2718**.

(238) Sub-decoder **3915** determines whether src1 bit field **1305** is 11110. If this is the case, sub-decoder **3915** supplies a stream 1 read signal to streaming engine **2700**. Streaming engine **2700** then supplies stream 1 data stored in holding register **2728** to the src1 input of functional unit **3920**.

(239) Sub-decoder **3916** determines whether src1 bit field **1305** is 11111. If this is the case, sub-decoder **3916** supplies a stream 1 read signal to streaming engine **2700**. Streaming engine **2700** then supplies stream 1 data stored in holding register **2728** to the src1 input of functional unit **3920**. Sub-decoder **3914** also supplies an advance signal to stream 1. As previously described, streaming engine **2700** advances to store the next sequential data elements of stream 1 in holding register **2728**.

(240) Similar circuits are used to select data supplied to scr2 input of functional unit **3902** in response to the bit coding of src2/cst field **1304**. The src2 input of functional unit **3920** may be supplied with a constant input in a manner described above.

(241) The exact number of instruction bits devoted to operand specification and the number of data registers and streams are design choices. Those skilled in the art would realize that other number selections that described in the application are feasible. In particular, the specification of a single global vector register file and omission of local vector register files is feasible. This invention employs a bit coding of an input operand selection field to designate a stream read and another bit coding to designate a stream read and advancing the stream.

(242) While this specification contains many specifics, these should not be construed as limitations on the scope of what may be claimed, but rather as descriptions of features that may be specific to particular embodiments. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment.

Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable sub-combination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a sub-combination or variation of a sub-combination.

(243) Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results unless such order is recited in one or more claims. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments.

Claims

1. A circuit device comprising: a processor that includes a register that includes a first portion and is configured to store a stream definition template that specifies: a set of loops; and whether the first portion of the register stores at least a portion of a measure of distance between starting points for consecutive iterations of a loop of the set of loops; and a circuit coupled to the processor that includes: an address generator configured to generate a set of addresses associated with a set of data using the stream definition template, wherein the address generator includes a counter circuit configured to: determine a current loop count of the loop; and multiply the current loop count by the measure of distance between starting points for consecutive iterations of the loop; and a

memory interface configured to couple between the address generator and a memory and configured to retrieve the set of data from the memory using the set of addresses.

2. The circuit device of claim 1, wherein: the register includes a second portion; and the stream definition template further specifies whether the second portion of the register stores at least a portion of a count of iterations of the loop of the set of loops.

3. The circuit device of claim 2, wherein: the loop is a first loop; the set of loops includes a second loop; and the stream definition template specifies whether the second portion of the register stores the at least a portion of the count of iterations of the first loop or at least a portion of a count of iterations of the second loop.

4. The circuit device of claim 1, wherein: the loop is a first loop; the set of loops includes a second loop; and the stream definition template specifies whether the second loop is enabled.

5. The circuit device of claim 4, wherein the second loop is an intermediate loop.

6. The circuit device of claim 4, wherein the second loop is an outer loop.

7. The circuit device of claim 1, wherein the stream definition template specifies whether the loop increases or decreases in address as the loop is traversed.

8. The circuit device of claim 1, wherein the stream definition template specifies whether each loop of the set of nested loops uses linear addressing or circular addressing.

9. The circuit device of claim 1 further comprising the memory coupled to the memory interface of the circuit, wherein the memory is a cache memory.

10. The circuit device of claim 9, wherein: the cache memory is a level two (L2) cache memory; the circuit device further comprises a level one (L1) cache memory coupled between the processor and the L2 cache memory; and the circuit is configured to couple between the processor and the L2 cache memory in parallel with the L1 cache memory.

11. The circuit device of claim 1, wherein: the circuit includes a set of multiplexers each associated with a respective loop of the set of nested loops; and the first portion of the register is coupled to an input of a first multiplexer of the set of multiplexers and to an input of a second multiplexer of the set of multiplexers.

12. A circuit device comprising: a circuit configured to receive a stream definition template and an instruction to read a set of data associated with the stream definition template, wherein: the stream definition template specifies: a set of loops associated with the set of data; and whether a first portion of the stream definition template specifies at least a portion of a measure of distance between starting points for consecutive iterations of a loop of the set of loops; and the circuit includes: an address generator configured to, based on the instruction, generate a set of addresses associated with the set of data using the stream definition template, wherein the address generator includes a counter circuit that includes a multiplier configured to multiply a count associated with the loop by the measure of distance between starting points for consecutive iterations of the loop; and a memory interface coupled to the address generator and configured to, based on the instruction, retrieve at least a portion of the set of data from a memory circuit using the set of addresses.

13. The circuit device of claim 12, wherein: the stream definition template includes a second portion; and the stream definition template further specifies whether a second portion of the stream definition template specifies at least a portion of a count of iterations of the loop of the set of loops.

14. The circuit device of claim 13, wherein: the loop is a first loop; the set of loops includes a second loop; and the stream definition template specifies whether the second portion of the stream definition template specifies at least a portion of the count of iterations of the first loop or at least a portion of a count of iterations of the second loop.

15. The circuit device of claim 12, wherein: the circuit includes a set of multiplexers each associated with a respective loop of the set of nested loops; the set of multiplexers includes a first multiplexer and a second multiplexer; and each of the first multiplexer and the second multiplexer are configured to receive the first portion of the stream definition template.

16. A device comprising: a processor that includes a register that includes a first portion and is configured to store a stream definition template that specifies: a set of loops; and whether the first portion of the register stores at least a portion of a measure of distance between starting points for consecutive iterations of a first loop of the set of loops; and a circuit coupled to the processor that includes: an address generator configured to generate a set of addresses associated with a set of data using the stream definition template, wherein: the address generator includes a set of counter circuits that includes a respective counter circuit for each loop of the set of loops; the set of counter circuits includes a first counter circuit for the first loop; and the first counter circuit is coupled to receive the measure of distance between starting points for consecutive iterations of the first loop; and a memory interface configured to couple between the address generator and a memory and configured to retrieve the set of data from the memory using the set of addresses.

17. The device of claim 16, wherein: the circuit includes a set of multiplexers that each include an output coupled to a respective counter circuit loop of the set of nested loops; the output of a first multiplexer of the set of multiplexers is coupled to the first counter circuit; the output of a second multiplexer of the set of multiplexers is coupled to a second counter circuit of the set of counter circuits; and the first portion of the register is coupled to an input of the first multiplexer and to an input of the second multiplexer.

18. The device of claim 16, wherein: the register includes a second portion; and the stream definition template further specifies whether the second portion of the register stores at least a portion of a count of iterations of the loop of the set of loops.

19. The device of claim 16 further comprising the memory coupled to the memory interface of the circuit, wherein the memory is a cache memory.

20. The device of claim 19, wherein: the cache memory is a level two (L2) cache memory; the device further comprises a level one (L1) cache memory coupled between the processor and the L2 cache memory; and the circuit is configured to couple between the processor and the L2 cache memory in parallel with the L1 cache memory.
