

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411775
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Cariello; Giuseppe et al.

Dual address encoding for logical-to-physical mapping

Abstract

Methods, systems, and devices for dual address encoding for logical-to-physical mapping are described. A memory device may identify a first physical address corresponding to a first logical block address generated by a host device and a second physical address corresponding to a second (consecutive) logical block address generated by a host device. The memory device may store the first physical address and second physical address in a single entry of a logical-to-physical mapping table that corresponds to the first logical block address. The memory device may transmit the logical-to-physical table to the host device for storage at the host device. The host device may subsequently transmit a single read command to the memory device that includes the first physical address and the second physical address based on the logical-to-physical table.

Inventors:	Cariello; Giuseppe (Boise, ID), Parry; Jonathan S. (Boise, ID)
Applicant:	Lodestar Licensing Group LLC (Evanston, IL)
Family ID:	78007808
Appl. No.:	18/662743
Filed:	May 13, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240296125 A1	Sep. 05, 2024

Related U.S. Application Data

continuation parent-doc US 18211476 20230619 US 12007903 child-doc US 18662743
continuation parent-doc US 17495410 20211006 US 11704252 20230718 child-doc US 18211476
continuation parent-doc US 16869397 20200507 US 11144471 20211012 child-doc US 17495410

Publication Classification

Int. Cl.: **G11C8/00** (20060101); **G06F12/10** (20160101); **G11C16/04** (20060101); **G11C16/10** (20060101); **G11C16/26** (20060101)

U.S. Cl.:

CPC **G06F12/10** (20130101); **G11C16/0483** (20130101); **G11C16/10** (20130101); **G11C16/26** (20130101); G06F2212/657 (20130101); G11C2216/14 (20130101)

Field of Classification Search

CPC: G11C (8/12); G11C (7/1021); G11C (16/10); G11C (7/1006); G11C (8/10)

USPC: 365/238.5

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10140215	12/2017	Canepa	N/A	G06F 3/064
10553278	12/2019	Narsale et al.	N/A	N/A
10733232	12/2019	Lee	N/A	N/A
11144471	12/2020	Cariello et al.	N/A	N/A
12007903	12/2023	Cariello	N/A	G11C 16/10
2011/0047347	12/2010	Li et al.	N/A	N/A
2018/0121121	12/2017	Mehra et al.	N/A	N/A
2019/0012227	12/2018	Kim et al.	N/A	N/A
2020/0004540	12/2019	Navon et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101030167	12/2006	CN	N/A
107066393	12/2016	CN	N/A
0090219	12/1982	EP	N/A

OTHER PUBLICATIONS

Chinese Patent Office, "Office Action and Search Report", issued in connection with Chinese Patent Application No. 202110473202.8, dated Mar. 23, 2022 (16 pages). cited by applicant

Primary Examiner: Dinh; Son T

Background/Summary

CROSS REFERENCE (1) This application is a continuation of U.S. patent application Ser. No. 18/211,476, filed Jun. 16, 2023, which is a continuation of U.S. patent application Ser. No. 17/495,410, filed Oct. 6, 2021, (U.S. Pat. No. 11,704,252), which is a continuation of U.S. patent

BACKGROUND

(1) The following relates generally to one or more memory systems and more specifically to dual address encoding for logical-to-physical mapping.

(2) Memory devices are widely used to store information in various electronic devices such as computers, wireless communication devices, cameras, digital displays, and the like. Information is stored by programing memory cells within a memory device to various states. For example, binary memory cells may be programmed to one of two supported states, often denoted by a logic 1 or a logic 0. To access the stored information, a component may read, or sense, at least one stored state in the memory device. To store information, a component may write, or program, the state in the memory device.

(3) Various types of memory devices exist, including magnetic hard disks, random access memory (RAM), read only memory (ROM), dynamic RAM (DRAM), synchronous dynamic RAM (SDRAM), ferroelectric RAM (FeRAM), magnetic RAM (MRAM), resistive RAM (RRAM), flash memory, phase change memory (PCM), 3-dimensional cross-point memory (3D Xpoint), Flash memory (such as floating-gate Flash and charge-trapping Flash, which may be used in not-or (NOR) or not-and (NAND) memory devices), and others. Memory devices may be volatile or non-volatile. Non-volatile memory cells, such as flash memory cells, may maintain their stored logic state for extended periods of time even in the absence of an external power source. Volatile memory cells, such as DRAM cells, may lose their stored state over time unless they are periodically refreshed by an external power source. Flash-based memory devices may have different performance compared to other non-volatile and volatile memory devices.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 illustrates an example of a memory device that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(2) FIG. 2 illustrates an example of a NAND circuit that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(3) FIG. 3 illustrates an example of a system that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(4) FIG. 4 illustrates an example of a set of entries that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(5) FIG. 5 illustrates an example of a process that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(6) FIG. 6 shows a block diagram of a memory device that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

(7) FIGS. 7 and 8 show flowcharts illustrating a method or methods that support dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein.

DETAILED DESCRIPTION

(8) Flash memory is generally organized into pages and blocks, with each block containing multiple pages. Flash memory cells may be read and written at the page level of granularity but may be erased at a block level of granularity. In some cases, Flash memory cells may be erased before they can be re-written with new data. Thus, when a Flash memory device updates a page of data (e.g., in response to a command from a host device), the memory device may write the new data to a different page and mark the old page as obsolete, rather than erasing a block of memory and re-writing the pages in the block.

(9) A memory device may receive commands, such as read and write commands for reading or writing data, from a host device. For a write operation, the host device may refer to the location of data stored in the memory device using a logical block address (LBA) that is mapped to a physical address of a page of memory of the memory device at which the data is stored. Because the physical address of the data may change (e.g., when data is updated by writing the updated data to a different page), some memory devices maintain one or more logical-to-physical (L2P) tables that map LBAs generated by the host device to corresponding physical addresses of pages in the memory device. In this manner, the host device can request to read data from the memory device using the same LBA even if the data has been moved to a different physical address of the memory device.

(10) In some cases, each entry in an L2P table may contain a single physical address that points to a page of data stored in the memory device. Because L2P entries may be frequently updated as data is written, over-written, moved, etc., L2P tables are often stored in DRAM or other memory component associated with the Flash memory, which supports relatively fast reads and writes.

(11) For memory devices with large memory capacities, however, L2P tables may be too big to store on the memory device itself. Moreover, some storage devices, such as Universal Flash Storage (UFS) devices, may lack on-chip DRAM for storing and updating L2P tables. Thus, some memory devices may use memory that resides on the host device (e.g., host DRAM) to store the L2P mapping tables rather than storing such tables locally. In such cases, during a read operation, the host device may generate the LBA and look up the corresponding physical address in the L2P table, then include the physical address in the read command to the memory device (e.g., rather than the memory device looking up the physical address based on an LBA received in a command from the host device).

(12) In some cases, each entry in an L2P table may include a single physical address that points to a single page of memory. Thus, reading more than one page of data may use multiple accesses of the L2P table to retrieve multiple physical addresses, thereby potentially increasing the read latency. In some cases, a host device may read multiple pages of logically consecutive data—that is, data that is stored for consecutive LBAs (but not necessarily at consecutive physical addresses). For example, analysis of cell phone data usage indicates that roughly 50% of random read traffic consists of single-page reads, with the remaining 50% consisting of multi-page reads.

(13) As described herein, to reduce the latency associated with reading multiple pages of data, the memory device may encode (e.g., pack, store) multiple physical addresses corresponding to multiple consecutive LBAs into a single entry of an L2P table. The host may retrieve such an entry of the L2P table based on the first LBA of the entry, and may include the multiple physical addresses (which may be non-consecutive physical addresses) of the single entry in a single read command that is sent to the memory device. The host device may also include an indication of a data transfer size (e.g., a quantity of pages or bytes to be read from the memory device) in the read command to indicate, to the memory device, how many physical addresses are included in the read command. Based on receiving the read command, the memory device may read the data from the multiple physical addresses and transmit the data to the host device.

(14) To enable access of each individual LBA and corresponding physical address in the L2P table, in some cases, a memory device may, when generating an L2P table, duplicate physical addresses in separate entries of the L2P table. For example, a first entry of an L2P table may include a first physical address corresponding to LBA 1 and a second physical address corresponding to LBA 2. A second entry of the L2P table may include (e.g., duplicate) the second physical address corresponding to LBA 2 and a third physical address corresponding to LBA 3. In this manner, the host device may be able to access each LBA individually, while also being able to retrieve two (or more) physical addresses from a single entry. In some cases, such an L2P table may also include L2P entries that contain a single physical address, thereby maintaining backwards compatibility with earlier systems.

(15) Features of the disclosure are initially described in the context of a memory device and NAND circuit as described with reference to FIGS. 1 and 2. Features of the disclosure are further described in the context of systems, sets of entries for dual-encoded L2P tables, and algorithms for generating dual-encoded L2P tables, as described with reference to FIGS. 3-5. These and other features of the disclosure are further illustrated by and described with reference to an apparatus diagram and flowcharts that relate to dual address encoding for L2P mapping as described with reference to FIGS. 6-8.

(16) FIG. 1 illustrates an example of a memory die **100** in accordance with examples as disclosed herein. In some cases, the memory die **100** may be referred to as a universal flash storage (UFS) device, a solid-state storage device, a managed memory device, a memory chip, or an electronic memory apparatus. The memory die **100** may include one or more memory cells, such as memory cell **105-a** and memory cell **105-b** (other memory cells are unlabeled). A memory cell **105** may be, for example, a Flash memory cell (such as depicted in the blow-up diagram of memory cell **105-a** shown in FIG. 1), a DRAM memory cell, an FeRAM memory cell, a PCM memory cell, or another type of memory cell.

(17) Each memory cell **105** may be programmed to store a logic state representing one or more bits of information. Different memory cell architectures may store a logic state in different ways. In FeRAM architectures, for example, each memory cell **105** may include a capacitor that includes a ferroelectric material to store a charge and/or a polarization representative of the programmable state. In DRAM architectures, each memory cell **105** may include a capacitor that includes a dielectric material (e.g., an insulator) to store a charge representative of the programmable state. In Flash memory architectures, each memory cell **105** may include a transistor that has a floating gate and/or a dielectric material for storing a charge representative of the logic state. For example, the blow-up diagram of memory cell **105-a** is a Flash memory cell that includes a transistor **110** (e.g., a metal-oxide-semiconductor (MOS) transistor) that may be used to store a logic state. The transistor **110** has a control gate **115** and may include a floating gate **120** that is sandwiched between dielectric material **125**. Transistor **110** includes a first node **130** (e.g., a source or drain) and a second node **135** (e.g., a drain or source). A logic state may be stored in transistor **110** by placing (e.g., writing, storing) a quantity of electrons (e.g., a charge) on floating gate **120**. The amount of charge to be stored on the floating gate **120** may depend on the logic state to be stored. The charge stored on floating gate **120** may affect the threshold voltage of transistor **110**, thereby affecting the amount of current that may flow through transistor **110** when transistor **110** is activated. The logic state stored in transistor **110** may be read by applying a voltage to the control gate **115** (e.g., at control node **140**) to activate transistor **110** and measuring (e.g., detecting, sensing) the resulting amount of current that flows between the first node **130** and the second node **135**.

(18) For example, a sense component **170** may determine a logic state stored on a Flash memory cell based on the presence or absence of a current from the memory cell, or based on whether the current is above or below a threshold current. Similarly, a Flash memory cell may be written by applying a voltage (e.g., a voltage above a threshold or a voltage below a threshold) to the memory cell to store (or not store) an electric charge on the floating gate representing one of the possible logic states.

(19) A charge-trapping Flash memory cell may operate in a manner similar to that of a floating-gate Flash memory cell, but instead of (or in addition to) storing a charge on a floating gate **120**, a charge-trapping Flash memory cell may store a charge representing the state in a dielectric material below the control gate **115**. Thus, a charge-trapping Flash memory cell may or may not include a floating gate **120**.

(20) In some examples, each row of memory cells **105** is connected to a word line **160** and each column of memory cells **105** is connected to a digit line **165**. Thus, one memory cell **105** may be located at the intersection of a word line **160** and a digit line **165**. This intersection may be referred to as a memory cell's address. Digit lines are sometimes referred to as bit lines. In some cases, word

lines **160** and digit lines **165** may be substantially perpendicular to one another and may create an array of memory cells **105** (e.g., in a memory array). In some cases, word lines **160** and digit lines **165** may be generically referred to as access lines or select lines.

(21) In some cases, memory die **100** may include a three-dimensional (3D) memory array, where multiple two-dimensional (1D) memory arrays are formed on top of one another. This may increase the quantity of memory cells that may be placed or created on a single die or substrate as compared with 1D arrays, which in turn may reduce production costs, or increase the performance of the memory array, or both. In the example of FIG. 1, memory die **100** includes multiple levels of memory arrays. The levels may, in some examples, be separated by an electrically insulating material. Each level may be aligned or positioned so that memory cells **105** may be aligned (exactly, overlapping, or approximately) with one another across each level, forming memory cell stack **175**. In some cases, memory cell stack **175** may be referred to as a string of memory cells, discussed in more detail with reference to FIG. 3.

(22) Accessing memory cells **105** may be controlled through row decoder **145** and column decoder **150**. For example, row decoder **145** may receive a row address from memory controller **155** and activate an appropriate word line **160** based on the received row address. Similarly, column decoder **150** may receive a column address from memory controller **155** and activate an appropriate digit line **165**. Thus, by activating one word line **160** and one digit line **165**, one memory cell **105** may be accessed.

(23) Upon accessing, memory cell **105** may be read, or sensed, by sense component **170**. For example, sense component **170** may be configured to determine the stored logic state of memory cell **105** based on a signal generated by accessing memory cell **105**. The signal may include a voltage or electrical current, or both, and sense component **170** may include voltage sense amplifiers, current sense amplifiers, or both. For example, a current or voltage may be applied to a memory cell **105** (using the corresponding word line **160** and/or digit line **165**) and the magnitude of the resulting current or voltage on the digit line **165** may depend on the logic state stored by the memory cell **105**. For example, for a Flash memory cell, the amount of charge stored on a floating gate or in an insulating layer of a transistor in the memory cell **105** may affect the threshold voltage of the transistor, thereby affecting the amount of current that flows through the transistor in the memory cell **105** when the memory cell **105** is accessed. Such differences in current may be used to determine the logic state stored on the memory cell **105**.

(24) Sense component **170** may include various transistors or amplifiers in order to detect and amplify a signal (e.g., a current or voltage) on a digit line **165**. The detected logic state of memory cell **105** may then be output via input/output block **180**. In some cases, sense component **170** may be a part of column decoder **150** or row decoder **145**, or sense component **170** may otherwise be connected to or in electronic communication with column decoder **150** or row decoder **145**.

(25) A memory cell **105** may be set or written by similarly activating the relevant word line **160** and digit line **165** to enable a logic state (e.g., representing one or more bits of information) to be stored in the memory cell **105**. Column decoder **150** or row decoder **145** may accept data, for example from input/output block **180**, to be written to the memory cells **105**. As previously discussed, in the case of Flash memory (such as Flash memory used in NAND and 3D NAND memory devices) a memory cell **105** may be written by storing electrons in a floating gate or an insulating layer.

(26) Memory controller **155** may control the operation (e.g., read, write, re-write, refresh, verify, erase) of memory cells **105** through the various components, for example, row decoder **145**, column decoder **150**, and sense component **170**. In some cases, one or more of row decoder **145**, column decoder **150**, and sense component **170** may be co-located with memory controller **155**. Memory controller **155** may generate row and column address signals in order to activate the desired word line **160** and digit line **165**. Memory controller **155** may also generate and control various voltages or currents used during the operation of memory die **100**.

(27) In some cases, memory controller **155** or another electronic component of memory die **100** may construct (e.g., build, generate, and/or maintain) one or more L2P tables for mapping LBAs generated by a host device to physical addresses in the memory die **100** (e.g., addresses of physical pages in memory die **100** that correspond to the LBAs). Memory die **100** may transmit such L2P tables to the host device, and the host device may store the L2P tables for subsequent look-up access for read operations.

(28) In some cases, memory die **100** may receive a read command from the host device that includes one or more physical addresses and an indication of a data transfer length. Memory die **100** may extract the physical addresses from the read command and may retrieve data from the multiple physical addresses based on the read command. Memory die **100** may transmit the retrieved data to the host device.

(29) FIG. 2 illustrates an example of NAND circuit **200** that supports dual address encoding for logical-to-physical mapping in accordance with examples of the present disclosure. NAND circuit **200** may be an example of a portion of a memory device, such as memory die **100**. Although some elements included in FIG. 2 are labeled with reference numbers, other corresponding elements are not labeled, though they are the same or would be understood to be similar, in an effort to increase visibility and clarity of the depicted features.

(30) NAND circuit **200** includes multiple Flash memory cells **205** (which may be, for example, Flash memory cells such as described with reference to FIG. 1) connected in a NAND configuration. In a NAND memory configuration (referred to as NAND memory), multiple Flash memory cells **205** are connected in series with each other to form strings **210** of memory cells **205**, in which the drain of each Flash memory cell **205** in the string **210** is coupled with the source of another Flash memory cell **205** in the string. In some cases, Flash memory cells that are connected in a NAND configuration to form a NAND memory may be referred to as NAND memory cells.

(31) Each string **210** of memory cells **205** may be associated with a corresponding digit line **215** (e.g., digit line **215-a**, **215-b**) that is shared by the memory cells **205** in the string **210**. Each memory cell **205** in a string **210** may be associated with a separate word line **230** (e.g., word line **230-a**, **230-i**, **230-n**), such that the quantity of word lines **230** may be equal to the quantity of memory cells **205** in a string **210**.

(32) NAND memory may be hierarchically organized as strings **210** that include multiple memory cells **205**, pages **255** that include memory cells **205** that are connected to the same word line **230** (e.g., memory cells **205** from multiple strings **210**), and blocks **260** that include multiple pages **255**. A NAND memory cell may be erased before it can be re-written. In some cases, NAND memory can be written to and read from at the page level of granularity (e.g., by activating the corresponding word line **230**), but may not be erasable at the page level of granularity. In some cases, NAND memory may instead be erasable at a higher level of granularity, such as at the block level of granularity. Different memory devices may have different read/write/erase characteristics.

(33) Each string **210** of memory cells **205** in NAND circuit **200** is coupled with a select gate device for drain (SGD) transistor **220** at one end of the string **210** and a select gate device for source (SGS) transistor **235** at the other end of the string **210**. SGD transistor **220** and SGS transistor **235** may be used to couple a string **210** of memory cells **205** to a digit line **215** and/or to a source node **250** (e.g., source node **250-a**, **250-b**) by applying a voltage at the gate **245** of SGD transistor **220** and/or at the gate **240** of SGS transistor **235**, respectively.

(34) During NAND memory operations, various voltage levels associated with source node **250**, gate **240** of an SGS transistor **235** associated with source node **250**, word lines **230**, drain node **225**, gate **245** of an SGD transistor **220** associated with drain node **225**, and digit line **215** may be applied to perform one or more operations (e.g., program, erase, or read) on at least some NAND memory cells in a string **210**.

(35) In some cases, during a read operation, a positive voltage may be applied to digit line **215** connected to drain node **225** whereas source node **250** may be connected to a ground or a virtual

ground (e.g., approximately 0 V). For example, the voltage applied to drain node **225** may be 1 V. Concurrently, voltages applied to gates **245** and **240** may be increased above the threshold voltages of the one or more SGSs **235** associated with source node **250** and the one or more SGD **220** associated with drain node **225**, such that a channel associated with string **210** may be electrically connected to drain node **225** and source node **250**. A channel may be an electrical path through the memory cells **205** in a string **210** (e.g., through the transistors in the memory cells **205**) that may conduct current under certain operating conditions.

(36) Concurrently, multiple word lines **230** (e.g., word lines **230-a**, **230-i**, **230-n**, or in some cases all word lines **230**) except a selected word line (i.e., word lines associated with unselected cells in string **210**) may be connected to a voltage (e.g., VREAD) that is higher than the highest threshold voltage (VT) of memory cells in string **210**. VREAD may cause some or all of the unselected memory cells in string **210** to turn “ON” so that each unselected memory cell can maintain high conductivity in a channel associated with it. In some examples, a word line **230** associated with a selected cell may be connected to a voltage, VTarget. VTarget may be selected at a value between VT of an erased memory cell and VT of a programmed memory cell in string **210**. When the selected memory cell exhibits an erased VT (e.g., VTarget > VT of the selected memory cell), the selected memory cell **205** may turn “ON” in response to the application of VTarget and thus allow a current to flow in the channel of string **210** from digit line **215** to source **250**. When the selected memory cell exhibits a programmed VT (e.g., hence VTarget < VT of the selected memory cell), the selected memory cell may turn “OFF” in response to VTarget and thus prohibit a current to flow in the channel of string **210** from digit line **215** to source **250**. The amount of current flow (or lack thereof), may be sensed by sense component **170** as described with reference to FIG. **1** to read stored information in the selected memory cell **205** within string **210**.

(37) FIG. **3** is an example of a system **300** that supports dual address encoding for logical-to-physical mapping in accordance with examples of the present disclosure. The system **300** includes a host device **305** coupled with a memory device **310**.

(38) Memory device **310** may be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of a storage device include a solid-state drive (SSD), a flash drive, a universal serial bus (USB) flash drive, an embedded Multi-Media Controller (eMMC) drive, a Universal Flash Storage (UFS) drive, and a hard disk drive (HDD). Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), and a non-volatile dual in-line memory module (NVDIMM).

(39) Host device **305** may use memory device **310** to store data in memory arrays **330** and read data from memory arrays **330**. Host device **305** may be a computing device such as a desktop computer, laptop computer, network server, mobile device, a vehicle (e.g., airplane, drone, train, automobile, or other conveyance), Internet of Things (IoT) enabled device, embedded computer (e.g., one included in a vehicle, industrial equipment, or a networked commercial device), or such computing device that includes a memory and a processing device. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, universal serial bus (USB) interface, Fibre Channel, Serial Attached SCSI (SAS), etc.

(40) Memory device **310** may include a memory device controller **325** and one or more memory die **330** for storing data. In some examples, memory device controller **325** may be included in a controller die that may be separate (e.g., distinct) from memory die **330** and may be packaged together with memory die **330** in a single package (e.g., a package containing memory device **310**).

(41) Memory die **330** may be an example of memory die **100** described with reference to FIG. **1** and may include NAND memory arrays, for example, or other types of memory arrays for reading and writing data for host device **305**.

(42) In some examples, each memory die **330** in memory device **310** may include a local memory controller, which may be responsible for orchestrating lower level operations, (such as activating

row and column drivers) during read operations, write operations, verify operations, erase operations, or other operations that may be performed on a memory array of memory die 330, as described with reference to memory controller 155 of FIG. 1.

(43) In some examples, memory device controller 325 may communicate with each of the local memory controllers (e.g., memory controller 155) of respective memory die 330 to direct memory die 330 to perform memory operations. Memory device controller 325 may also handle some higher-level operations of memory device 310, such as garbage collection or other operations. In some examples, the term “controller” as used herein may refer to memory device controller 325, a local memory controller 155, or a combination thereof.

(44) In some cases, memory device 310 may also include static random access memory (SRAM) memory 340 or other types of memory that may be used by memory device 310 for internal storage or calculations, for example. Accesses to SRAM memory 340 may be faster than accesses to a memory array (e.g., a NAND memory array) in memory die 330, and therefore it may be desirable to use SRAM memory 340 for storing and updating L2P tables. In some examples, however, SRAM memory 340 may be integrated with memory device controller 325 on a single die, which may constrain the size of SRAM memory 340 (e.g., due to cost or other constraints associated with memory device controller 325). Thus, in some examples, a size of SRAM memory 340 may be smaller than a size of an L2P table used by memory device 310 and may therefore be insufficient for storing L2P mapping information. As a result, in some examples, a memory device 310 may generate L2P mapping tables and transmit them to host device 305 for storage rather than storing the L2P tables locally on memory die 330. As described in more detail with reference to FIG. 5, in some examples, SRAM memory 340 may be used, by memory device 310, for building one or more L2P tables that may be transmitted to host device 305 for storage.

(45) Host device 305 includes host controller interface 320. Host controller interface 320 may provide an interface for passing control, address, data, and other signals between host device 305 and memory device 310. Host device 305 may transmit memory access commands, such as read or write commands, to memory device 310 using host controller interface 320.

(46) Memory device controller 325 may receive signals from host device 305 via host controller interface 320 and may cause memory device 310 to perform certain operations in response to receiving such signals. For example, memory device controller 325 may receive a read or write command from host device 305 and, in response, may cause memory device 310 to read data or write data to memory die 330 based on the received command.

(47) Host device 305 includes host memory 315, which may include one or more types of volatile or non-volatile memory. For example, host memory 315 may include SRAM, DRAM, Flash, or other types of memory.

(48) In some examples, memory device 310 may build and maintain one or more sets of entries (e.g., L2P look-up tables) for mapping LBAs generated by host device 305 to physical addresses (e.g., page addresses) of memory die 330. Such sets of entries may be generated based on receiving one or more write commands from the host device 305 that include an LBA for writing data. Each entry in an L2P table may include one or more physical addresses corresponding to the one or more LBAs in the write command(s). In some cases, memory device 310 may transmit such sets of entries to host device 305 such that the sets of entries (L2P tables) are stored in host memory 315. For example, memory device controller 325 may transmit such sets of entries to host device 305 (e.g., via host controller interface 320).

(49) The memory device 310 may include an accelerator 345 for building the sets of entries. In some cases, some or all of accelerator 345 may be a hardware accelerator configured to perform a dedicated function in memory device 310. In some cases, some or all of accelerator 345 may be implemented in hardware, software, firmware, or a combination thereof that is executed by memory device controller 325, or by a local memory controller of memory die 330, or by another component of memory device 310, or by some combination of these.

(50) In some examples, entries in a set of entries may be ordered sequentially by an LBA index. That is, a first entry in a set of entries may correspond to LBA **0**, a second (consecutive) entry in the set of entries may correspond to LBA **1**, a third entry may correspond to LBA **2**, and so on. In some examples, memory device **310** may encode (e.g., pack, store) two or more physical addresses (e.g., corresponding to two or more consecutive LBAs) into a single entry of the set of entries. For example, if memory device **310** receives multiple write commands from host device **305** that specify consecutive (e.g., LBAs having indices that are numerically consecutive, contiguous) LBAs, memory device **310** may determine (e.g., select, identify) multiple corresponding physical addresses at which the data will be written, and may store, in the set of entries, a single entry that contains the multiple physical addresses.

(51) When host device **305** reads data from memory device **310**, host device **305** may identify one or more LBAs of the data to be read, and may look up corresponding physical addresses in the set of entries that resides in host memory **315**. Host device **305** may, based on looking up an entry in the set of entries, transmit a read command that includes the physical address(es), the LBA, and an indication of a data transfer length to memory device **310** (e.g., whether the read operation is a 4 kB read or an 8 kB read). Memory device **310** may then retrieve data at the physical address(es) specified in the read command and transmit the retrieved data to host device **305**. For example, host device **305** may transmit a read command that includes a first physical address corresponding to a first LBA and a second physical address corresponding to a second (consecutive) LBA. The read command may include an indication that the data transfer length is two (e.g., two pages); e.g., that the read command includes two physical addresses for a data transfer length of two pages. The memory device **310** may retrieve a first portion of the data from the first physical address and a second portion of the data from the second physical address and transmit the first portion of the data and second portion of the data to the host device **305**. The combined length of the first portion of the data and the second portion of the data may be associated with (e.g., equal to, corresponding to) the indicated data transfer length. For example, if the indicated data transfer length is two, the combined length of the first portion of the data and the second portion of the data may be two pages (e.g., 8 kB, for a memory device with 4 kB pages).

(52) FIG. 4 illustrates an example of a set of entries **400** that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein. Set of entries **400** may represent an L2P table that may be generated by a memory device (such as memory device **310**) for mapping LBAs to physical addresses. In some cases, set of entries **400** may be generated by a memory device, transmitted to a host device such as host device **305** described with reference to FIG. 3, and stored in memory that resides on the host device.

(53) In some cases, a memory device may generate set of entries **400** based on receiving multiple write commands from a host device that specify multiple consecutive LBAs (e.g., LBA **1**, LBA **2**, LBA **3**, etc.) for writing data to the memory device. Based on receiving the write commands, the memory device may select physical addresses for writing the data, write the data to the physical addresses, and generate one or more entries in set of entries **400** that map the LBAs generated by the host device to the corresponding physical addresses at which the data is written. Entries in the set of entries **400** may be ordered by the index of the LBAs; e.g., a first entry associated with LBA **0**, a second entry associated with LBA **1**, etc. The physical addresses corresponding to the LBAs and stored in the set of entries **400** may not be consecutive. In some cases, a memory device may include an accelerator for building a set of entries **400** more efficiently, as described with reference to FIG. 5.

(54) In some cases, each entry in set of entries **400** consumes eight (8) bytes, and may map each LBA to a page of physical addresses on the memory device. A page of memory may be, for example, 4 kB of memory, 8 kB, or another size. For 4 kB pages, a set of entries that includes 1024 entries (such as set of entries **400**) may be used to map LBAs to physical addresses for up to 4 MB of memory. In some cases, each entry of set of entries **400** may be accessed (e.g., by a host device)

based on an LBA index of a first portion of each entry.

(55) Set of entries **400** includes first entry **410** and second entry **415**, among other entries. First entry **410** includes a first portion **410-a** that includes a first physical address corresponding to LBA **0** and a second portion **410-b** that includes a second physical address corresponding to LBA **1**. As previously noted, the first physical address and second physical address may not be consecutive, although they may be consecutive in some cases.

(56) Second entry **415** includes a first portion **415-a** that includes the second physical address corresponding to LBA **1** and a second portion **415-b** that includes a third physical address corresponding to LBA **2**. Thus, the physical address corresponding to LBA **1** is duplicated in second portion **410-b** and first portion **415-a**. More broadly, the physical addresses corresponding to LBAs **1-1023** are duplicated across entries to enable individual retrieval of each physical address based on the LBA index of the first portion of the entry.

(57) In some cases, set of entries **400** may include entries that contain a single physical address. For example, a final entry **420** of set of entries **400** may include a single physical address, and there may be other entries in set of entries **400** that include a single physical address.

(58) In some cases, an entry in the set of entries may encode additional information (e.g., in addition to the one or more physical addresses). For example, an entry may include data verification information for the one or more physical addresses. For example, when the memory device generates the set of entries, the memory device may determine an error detection value (e.g., a checksum or exclusive-or (XOR) value) for the physical address(es) and store the error detection value in the entry with the physical address(es).

(59) Although set of entries **400** depicts each entry (e.g., entries **410**, **415**) as including two physical addresses (corresponding to two consecutive LBAs), in some cases, a set of entries may include three or more physical addresses (corresponding to three or more consecutive LBAs).

(60) Moreover, although set of entries **400** includes 1024 eight-byte entries that are each split into two four-byte portions, other quantities of entries, entry lengths, and/or portion lengths may be used without departing from the scope of the disclosure.

(61) The implementation of set of entries **400** depicted in FIG. **4** may enable a host device to read data for multiple LBAs based on retrieving a single entry from set of entries **400**, thereby reducing the overhead and latency associated with reading multiple blocks of data.

(62) For example, if a host device intends to read data corresponding to LBA **1** and LBA **2**, the host device may retrieve second entry **415**. The host device may include the two physical addresses (e.g., a first physical address in first portion **415-a** and a second physical address in second portion **415-b**) in the read command, along with the first LBA (LBA**1**). The host device may also include, in the read command, an indication that the data transfer length is 2 (corresponding to reading data at two physical addresses). For example, a 4 kB read operation may correspond to reading a single physical address associated with a logical block address, whereas an 8 kB read operation may corresponds to reading two physical addresses with the logical block address.

(63) A memory device may receive the read command, and may extract the first physical address (e.g., in first portion **415-a**), the second physical address (e.g., in second portion **415-b**), and the indication that the data transfer length is 2. The memory device may, based on the read command and the indication of the data transfer length, retrieve a first portion of the data that is stored at the first physical address and a second portion of the data that is stored at the second physical address, and may transmit the first portion of the data and the second portion of the data to the host device.

(64) FIG. **5** illustrates an example of a process **500** that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein. Process **500** may be used to efficiently build a set of entries (such as set of entries **400**) for mapping LBAs to physical addresses by encoding multiple physical addresses in a single L2P entry. That is, process **500** may be used to convert an L2P table that has one physical address per entry to an L2P table that has two physical addresses per entry, as depicted in FIG. **4**.

(65) In some cases, process **500** may be implemented using a hardwired hardware accelerator, such as accelerator **345** depicted in FIG. **3**. In some cases, process **500** may be implemented using firmware or software. In the description below, process **500** is described as being implemented by an accelerator; such description should not be interpreted as limiting.

(66) At **505**, an accelerator of a memory device (such as accelerator **345**) may begin to build a set of entries (e.g., an L2P table) that encodes multiple physical addresses in one or more entries of the set of entries for mapping LBAs generated by a host device to physical addresses of a memory device.

(67) At **510**, the accelerator may initialize a first pointer to a first value, where the first pointer points to a location of a source memory area. The accelerator may initialize a second pointer to a first value, where the second pointer points to a location of a target memory area. In some cases, the source memory area may contain a previously generated L2P table having a single physical address in each entry. The source memory area may include one zone of SRAM memory on a memory device (such as one zone of SRAM **340**) and the target memory area may include another zone of SRAM memory on the memory device (such as another zone of SRAM **340**).

(68) At **515**, the accelerator may read the value of the first pointer S (e.g., a first physical address) and save this value in a first portion (**P0**) of a first entry of a set of entries, such as first portion **410-a** described with reference to FIG. **4**. The accelerator may increase the value of the first pointer S from the first value to a second value, such as by incrementing the first pointer S.

(69) At **520**, the accelerator may read the (incremented) value of the first pointer S (e.g., a second physical address) and save this value in a second portion (**P1**) of an entry of a set of entries, such as second portion **410-b** described with reference to FIG. **4**. The accelerator may again increment the first pointer S.

(70) At **525**, the accelerator may write the value of the first portion (**P0**) of the first entry at a memory location in the target memory indicated by second pointer T. The accelerator may increase the value of the second pointer T from the first value to a second value, such as by incrementing the second pointer T.

(71) At **530**, the accelerator may write the value of the second portion (**P1**) of the first entry at a memory location in the target memory indicated by the (incremented) second pointer T. The accelerator may again increment the second pointer T.

(72) At **535**, the accelerator may determine whether some or all of the entries stored in the source memory have been read into **P0** or **P1**; e.g., whether the original L2P table has been fully converted to a dual-encoded L2P table.

(73) In response to determining that all of the entries have been read into **P0** or **P1**, at **540**, the accelerator may cease processing. In some cases, the accelerator may cease processing based on the quantity of entries created satisfying a threshold. The memory device may then transmit the dual-encoded set of entries to a host device for storing in memory on the host device.

(74) In response to determining that some or all of the entries have not been read into **P0** or **P1**, at **545**, the accelerator may write the value of the second portion (**P1**) of the entry to the first portion (**P0**) of the next entry (e.g., duplicating the physical address in the first portion of the next entry of the L2P table), and continue to the next entry of the source memory table. In some cases, the accelerator may continue processing based on the quantity of entries created not satisfying a threshold.

(75) FIG. **6** shows a block diagram **600** of a memory device **605** that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein. The memory device **605** may be an example of aspects of a memory device as described with reference to FIGS. **1** through **5**. The memory device **605** may include a command receiving component **610**, a data transmission component **615**, a data writing component **620**, an entry generation component **625**, an LBA identification component **630**, a table transmission component **635**, and a data reading component **640**. Each of these modules may communicate, directly or indirectly, with one another

(e.g., via one or more buses).

(76) The command receiving component **610** may receive, from a host device, a read command including a first physical address of the memory device corresponding to a first logical block address generated by the host device and a second physical address of the memory device corresponding to a second logical block address generated by the host device. In some cases, the read command includes an indication of a data transfer length, and where a combined length of the first page of data and the second page of data is associated with the data transfer length. In some cases, the first physical address and the second physical address are contained within 8 bytes of the read command. In some cases, the read command includes the first logical block address. In some cases, the first logical block address and the second logical block address are consecutive logical block addresses. In some cases, the first physical address and the second physical address are non-consecutive physical addresses.

(77) In some examples, the command receiving component **610** may receive, at the memory device from the host device before receiving the read command, a write command including the first logical block address, where the write command is associated with writing the first page of data to the memory device.

(78) The data transmission component **615** may transmit the first page of data and the second page of data to the host device based on retrieving the first page of data and the second page of data.

(79) The data transmission component **615** may transmit, to the host device before receiving the read command, a set of entries that maps a set of logical block addresses to a corresponding set of physical addresses, where a first entry in the set of entries includes the first physical address and the second physical address, and a second entry includes the second physical address and a third physical address associated with a third logical block address.

(80) The entry generation component **625** may store a first entry of a set of entries, the first entry including the first physical address and the second physical address, the first entry for mapping the first logical block address to the first physical address and the second physical address. In some cases, each entry of the set of entries has a length of 8 bytes.

(81) In some examples, the entry generation component **625** may store a second entry of the set of entries, the second entry including the second physical address and a third physical address associated with a third logical block address, the second entry for mapping the second logical block address to the second physical address and third physical address. In some examples, the entry generation component **625** may store, at the memory device, a first entry in a set of entries based on writing the first page of data to the memory device at the first physical address and the second page of data to the memory device at the second physical address.

(82) In some examples, the entry generation component **625** may retrieve the first physical address from a first memory area based on a value of a first pointer. In some examples, the entry generation component **625** may store the first physical address in a first portion of the first entry.

(83) In some examples, the entry generation component **625** may increase the value of the first pointer to a second value of the first pointer. In some examples, the entry generation component **625** may retrieve the second physical address from the first memory area based on the second value of the first pointer. In some examples, the entry generation component **625** may store the second physical address in a second portion of the first entry.

(84) In some examples, the entry generation component **625** may store the first physical address in a second memory area based on a first value of a second pointer. In some examples, the entry generation component **625** may increase the first value of the second pointer to a second value of the second pointer.

(85) In some examples, the entry generation component **625** may store the second physical address in the second memory area based on the second value of the second pointer, where storing the second entry includes storing the second physical address in the first portion of the second entry. In some examples, the entry generation component **625** may determine an error detection value

associated with the first physical address and the second physical address, where the first entry includes the error detection value.

(86) The LBA identification component **630** may identify a first logical block address associated with a first physical address of the memory device. In some examples, the LBA identification component **630** may identify a second logical block address associated with a second physical address of the memory device.

(87) The table transmission component **635** may transmit the set of entries to a host device.

(88) The data writing component **620** may write the first page of data to the memory device at the first physical address associated with the first logical block address based on the write command. In some examples, the data writing component **620** may write the second page of data to the memory device at the second physical address based on a second write command.

(89) The data reading component **640** may retrieve a first page of data from the first physical address of the memory device and a second page of data from the second physical address of the memory device.

(90) FIG. 7 shows a flowchart illustrating a method or methods **700** that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein. The operations of method **700** may be implemented by a memory device or its components as described herein. For example, the operations of method **700** may be performed by a memory device as described with reference to FIGS. 1 through 6. In some examples, a memory device may execute a set of instructions to control the functional elements of the memory device to perform the described functions. Additionally or alternatively, a memory device may perform aspects of the described functions using special-purpose hardware.

(91) At **705**, the memory device may receive, from a host device, a read command including a first physical address of the memory device corresponding to a first logical block address generated by the host device and a second physical address of the memory device corresponding to a second logical block address generated by the host device. The operations of **705** may be performed according to the methods described herein. In some examples, aspects of the operations of **705** may be performed by a command receiving component as described with reference to FIG. 6.

(92) At **710**, the memory device may retrieve a first page of data from the first physical address of the memory device and a second page of data from the second physical address of the memory device. The operations of **710** may be performed according to the methods described herein. In some examples, aspects of the operations of **710** may be performed by a data reading component as described with reference to FIG. 6.

(93) At **715**, the memory device may transmit the first page of data and the second page of data to the host device based on retrieving the first page of data and the second page of data. The operations of **715** may be performed according to the methods described herein. In some examples, aspects of the operations of **715** may be performed by a data transmission component as described with reference to FIG. 6.

(94) In some examples, an apparatus as described herein may perform a method or methods, such as the method **700**. The apparatus may include features, means, or instructions (e.g., a non-transitory computer-readable medium storing instructions executable by a processor) for receiving, from a host device, a read command including a first physical address of the memory device corresponding to a first logical block address generated by the host device and a second physical address of the memory device corresponding to a second logical block address generated by the host device, retrieving a first page of data from the first physical address of the memory device and a second page of data from the second physical address of the memory device, and transmitting the first page of data and the second page of data to the host device based on retrieving the first page of data and the second page of data.

(95) In some examples of the method **700** and the apparatus described herein, the read command includes an indication of a data transfer length, and a combined length of the first page of data and

the second page of data may be associated with the data transfer length.

(96) Some examples of the method **700** and the apparatus described herein may further include operations, features, means, or instructions for transmitting, to the host device before receiving the read command, a set of entries that maps a set of logical block addresses to a corresponding set of physical addresses, where a first entry in the set of entries includes the first physical address and the second physical address, and a second entry includes the second physical address and a third physical address associated with a third logical block address.

(97) In some examples of the method **700** and the apparatus described herein, each entry of the set of entries may have a length of 8 bytes.

(98) Some examples of the method **700** and the apparatus described herein may further include operations, features, means, or instructions for receiving, at the memory device from the host device before receiving the read command, a write command including the first logical block address, where the write command may be associated with writing the first page of data to the memory device, and writing the first page of data to the memory device at the first physical address associated with the first logical block address based on the write command.

(99) Some examples of the method **700** and the apparatus described herein may further include operations, features, means, or instructions for storing, at the memory device, a first entry in a set of entries based on writing the first page of data to the memory device at the first physical address and writing the second page of data to the memory device at the second physical address.

(100) In some examples of the method **700** and the apparatus described herein, the first physical address and the second physical address may be contained within 8 bytes of the read command.

(101) In some examples of the method **700** and the apparatus described herein, the read command includes the first logical block address.

(102) In some examples of the method **700** and the apparatus described herein, the first logical block address and the second logical block address may be consecutive logical block addresses.

(103) In some examples of the method **700** and the apparatus described herein, the first physical address and the second physical address may be non-consecutive physical addresses.

(104) FIG. **8** shows a flowchart illustrating a method or methods **800** that supports dual address encoding for logical-to-physical mapping in accordance with examples as disclosed herein. The operations of method **800** may be implemented by a memory device or its components as described herein. For example, the operations of method **800** may be performed by a memory device as described with reference to FIG. **6**. In some examples, a memory device may execute a set of instructions to control the functional elements of the memory device to perform the described functions. Additionally or alternatively, a memory device may perform aspects of the described functions using special-purpose hardware, such as a hardware accelerator.

(105) At **805**, the memory device may identify a first logical block address associated with a first physical address of the memory device. The operations of **805** may be performed according to the methods described herein. In some examples, aspects of the operations of **805** may be performed by an LBA identification component as described with reference to FIG. **6**.

(106) At **810**, the memory device may identify a second logical block address associated with a second physical address of the memory device. The operations of **810** may be performed according to the methods described herein. In some examples, aspects of the operations of **810** may be performed by an LBA identification component as described with reference to FIG. **6**.

(107) At **815**, the memory device may store a first entry of a set of entries, the first entry including the first physical address and the second physical address, the first entry for mapping the first logical block address to the first physical address and the second physical address. The operations of **815** may be performed according to the methods described herein. In some examples, aspects of the operations of **815** may be performed by an entry generation component as described with reference to FIG. **6**.

(108) At **820**, the memory device may store a second entry of the set of entries, the second entry

including the second physical address and a third physical address associated with a third logical block address, the second entry for mapping the second logical block address to the second physical address and third physical address. The operations of **820** may be performed according to the methods described herein. In some examples, aspects of the operations of **820** may be performed by an entry generation component as described with reference to FIG. 6.

(109) At **825**, the memory device may transmit the set of entries to a host device. The operations of **825** may be performed according to the methods described herein. In some examples, aspects of the operations of **825** may be performed by a table transmission component as described with reference to FIG. 6.

(110) In some examples, an apparatus as described herein may perform a method or methods, such as the method **800**. The apparatus may include features, means, or instructions (e.g., a non-transitory computer-readable medium storing instructions executable by a processor) for identifying a first logical block address associated with a first physical address of the memory device, identifying a second logical block address associated with a second physical address of the memory device, storing a first entry of a set of entries, the first entry including the first physical address and the second physical address, the first entry for mapping the first logical block address to the first physical address and the second physical address, storing a second entry of the set of entries, the second entry including the second physical address and a third physical address associated with a third logical block address, the second entry for mapping the second logical block address to the second physical address and the third physical address, and transmitting the set of entries to a host device.

(111) In some examples of the method **800** and the apparatus described herein, storing the first entry may include operations, features, means, or instructions for retrieving the first physical address from a first memory area based on a value of a first pointer, storing the first physical address in a first portion of the first entry, increasing the value of the first pointer to a second value of the first pointer, retrieving the second physical address from the first memory area based on the second value of the first pointer, and storing the second physical address in a second portion of the first entry.

(112) Some examples of the method **800** and the apparatus described herein may further include operations, features, means, or instructions for storing the first physical address in a second memory area based on a first value of a second pointer, increasing the first value of the second pointer to a second value of the second pointer, and storing the second physical address in the second memory area based on the second value of the second pointer, where storing the second entry includes storing the second physical address in the first portion of the second entry.

(113) In some examples of the method **800** and the apparatus described herein, the first portion of the first entry and the second portion of the first entry each may have a size of four bytes.

(114) Some examples of the method **800** and the apparatus described herein may further include operations, features, means, or instructions for determining an error detection value associated with the first physical address and the second physical address, where the first entry include the error detection value.

(115) It should be noted that the methods described herein are possible implementations, and that the operations and the steps may be rearranged or otherwise modified and that other implementations are possible. Furthermore, portions from two or more of the methods may be combined.

(116) Information and signals described herein may be represented using any of a variety of different technologies and techniques. For example, data, instructions, commands, information, signals, bits, symbols, and chips that may be referenced throughout the above description may be represented by voltages, currents, electromagnetic waves, magnetic fields or particles, optical fields or particles, or any combination thereof. Some drawings may illustrate signals as a single signal; however, it will be understood by a person of ordinary skill in the art that the signal may represent

a bus of signals, where the bus may have a variety of bit widths.

(117) The terms “electronic communication,” “conductive contact,” “connected,” and “coupled” may refer to a relationship between components that supports the flow of signals between the components. Components are considered in electronic communication with (or in conductive contact with or connected with or coupled with) one another if there is any conductive path between the components that can, at any time, support the flow of signals between the components. At any given time, the conductive path between components that are in electronic communication with each other (or in conductive contact with or connected with or coupled with) may be an open circuit or a closed circuit based on the operation of the device that includes the connected components. The conductive path between connected components may be a direct conductive path between the components or the conductive path between connected components may be an indirect conductive path that may include intermediate components, such as switches, transistors, or other components. In some examples, the flow of signals between the connected components may be interrupted for a time, for example, using one or more intermediate components such as switches or transistors.

(118) The term “coupling” refers to condition of moving from an open-circuit relationship between components in which signals are not presently capable of being communicated between the components over a conductive path to a closed-circuit relationship between components in which signals can be communicated between components over the conductive path. When a component, such as a controller, couples other components together, the component initiates a change that allows signals to flow between the other components over a conductive path that previously did not permit signals to flow.

(119) The devices discussed herein, including a memory array, may be formed on a semiconductor substrate, such as silicon, germanium, silicon-germanium alloy, gallium arsenide, gallium nitride, etc. In some examples, the substrate is a semiconductor wafer. In other cases, the substrate may be a silicon-on-insulator (SOI) substrate, such as silicon-on-glass (SOG) or silicon-on-sapphire (SOS), or epitaxial layers of semiconductor materials on another substrate. The conductivity of the substrate, or sub-regions of the substrate, may be controlled through doping using various chemical species including, but not limited to, phosphorous, boron, or arsenic. Doping may be performed during the initial formation or growth of the substrate, by ion-implantation, or by any other doping means.

(120) A switching component or a transistor discussed herein may represent a field-effect transistor (FET) and comprise a three terminal device including a source, drain, and gate. The terminals may be connected to other electronic elements through conductive materials, e.g., metals. The source and drain may be conductive and may comprise a heavily-doped, e.g., degenerate, semiconductor region. The source and drain may be separated by a lightly-doped semiconductor region or channel. If the channel is n-type (i.e., majority carriers are electrons), then the FET may be referred to as a n-type FET. If the channel is p-type (i.e., majority carriers are holes), then the FET may be referred to as a p-type FET. The channel may be capped by an insulating gate oxide. The channel conductivity may be controlled by applying a voltage to the gate. For example, applying a positive voltage or negative voltage to an n-type FET or a p-type FET, respectively, may result in the channel becoming conductive. A transistor may be “on” or “activated” when a voltage greater than or equal to the transistor's threshold voltage is applied to the transistor gate. The transistor may be “off” or “deactivated” when a voltage less than the transistor's threshold voltage is applied to the transistor gate.

(121) The description set forth herein, in connection with the appended drawings, describes example configurations and does not represent all the examples that may be implemented or that are within the scope of the claims. The term “exemplary” used herein means “serving as an example, instance, or illustration,” and not “preferred” or “advantageous over other examples.” The detailed description includes specific details to providing an understanding of the described

techniques. These techniques, however, may be practiced without these specific details. In some instances, well-known structures and devices are shown in block diagram form to avoid obscuring the concepts of the described examples.

(122) In the appended figures, similar components or features may have the same reference label. Further, various components of the same type may be distinguished by following the reference label by a dash and a second label that distinguishes among the similar components. If just the first reference label is used in the specification, the description is applicable to any one of the similar components having the same first reference label irrespective of the second reference label.

(123) Information and signals described herein may be represented using any of a variety of different technologies and techniques. For example, data, instructions, commands, information, signals, bits, symbols, and chips that may be referenced throughout the above description may be represented by voltages, currents, electromagnetic waves, magnetic fields or particles, optical fields or particles, or any combination thereof.

(124) The various illustrative blocks and modules described in connection with the disclosure herein may be implemented or performed with a general-purpose processor, a digital signal processor (DSP), an application-specific integrated circuit (ASIC), a field-programmable gate array (FPGA) or other programmable logic device, discrete gate or transistor logic, discrete hardware components, or any combination thereof designed to perform the functions described herein. A general-purpose processor may be a microprocessor, but in the alternative, the processor may be any processor, controller, microcontroller, or state machine. A processor may also be implemented as a combination of computing devices (e.g., a combination of a DSP and a microprocessor, multiple microprocessors, one or more microprocessors in conjunction with a DSP core, or any other such configuration).

(125) The functions described herein may be implemented in hardware, software executed by a processor, firmware, or any combination thereof. If implemented in software executed by a processor, the functions may be stored on or transmitted over as one or more instructions or code on a computer-readable medium. Other examples and implementations are within the scope of the disclosure and appended claims. For example, due to the nature of software, functions described above can be implemented using software executed by a processor, hardware, firmware, hardwiring, or combinations of any of these. Features implementing functions may also be physically located at various positions, including being distributed such that portions of functions are implemented at different physical locations. Also, as used herein, including in the claims, “or” as used in a list of items (for example, a list of items prefaced by a phrase such as “at least one of” or “one or more of”) indicates an inclusive list such that, for example, a list of at least one of A, B, or C means A or B or C or AB or AC or BC or ABC (i.e., A and B and C). Also, as used herein, the phrase “based on” shall not be construed as a reference to a closed set of conditions. For example, an exemplary step that is described as “based on condition A” may be based on both a condition A and a condition B without departing from the scope of the present disclosure. In other words, as used herein, the phrase “based on” shall be construed in the same manner as the phrase “based at least in part on.”

(126) Computer-readable media includes both non-transitory computer storage media and communication media including any medium that facilitates transfer of a computer program from one place to another. A non-transitory storage medium may be any available medium that can be accessed by a general purpose or special purpose computer. By way of example, and not limitation, non-transitory computer-readable media can comprise RAM, ROM, electrically erasable programmable read-only memory (EEPROM), compact disk (CD) ROM or other optical disk storage, magnetic disk storage or other magnetic storage devices, or any other non-transitory medium that can be used to carry or store desired program code means in the form of instructions or data structures and that can be accessed by a general-purpose or special-purpose computer, or a general-purpose or special-purpose processor. Also, any connection is properly termed a computer-

readable medium. For example, if the software is transmitted from a website, server, or other remote source using a coaxial cable, fiber optic cable, twisted pair, digital subscriber line (DSL), or wireless technologies such as infrared, radio, and microwave, then the coaxial cable, fiber optic cable, twisted pair, digital subscriber line (DSL), or wireless technologies such as infrared, radio, and microwave are included in the definition of medium. Disk and disc, as used herein, include CD, laser disc, optical disc, digital versatile disc (DVD), floppy disk and Blu-ray disc where disks usually reproduce data magnetically, while discs reproduce data optically with lasers.

Combinations of the above are also included within the scope of computer-readable media.

(127) The description herein is provided to enable a person skilled in the art to make or use the disclosure. Various modifications to the disclosure will be apparent to those skilled in the art, and the generic principles defined herein may be applied to other variations without departing from the scope of the disclosure. Thus, the disclosure is not limited to the examples and designs described herein but is to be accorded the broadest scope consistent with the principles and novel features disclosed herein.

Claims

1. A method, comprising: transmitting a read command comprising: a transfer length field indicating a quantity of two or more contiguous logical blocks of data to be returned; and a value associated with a set of entries that comprise respective physical block addresses of the two or more contiguous logical blocks of data; and receiving the two or more contiguous logical blocks of data read from the respective physical block addresses based at least in part on the transfer length field and the set of entries.
2. The method of claim 1, wherein the read command further comprises a logical block address of a first logical block of the two or more contiguous logical blocks of data and an entry containing a physical block address of the first logical block.
3. The method of claim 1, further comprising: transmitting the set of entries comprising the respective physical block addresses of the two or more contiguous logical blocks of data.
4. The method of claim 1, wherein a size of each of the two or more contiguous logical blocks of data is 4 kilobytes (kB).
5. The method of claim 1, wherein a value of the transfer length field indicates a combined size of the two or more contiguous logical blocks of data.
6. The method of claim 1, wherein each entry of the set of entries is maintained in a cache of a host device.
7. The method of claim 1, wherein each entry of the set of entries has a length of 8 bytes.
8. A method, comprising: receiving a read command comprising: a transfer length field indicating a quantity of two or more contiguous logical blocks of data to be returned; and a value associated with a set of entries that comprise respective physical block addresses of the two or more contiguous logical blocks of data; storing the set of entries in association with the value; and transmitting the two or more contiguous logical blocks of data read from the respective physical block addresses based at least in part on the transfer length field and the set of entries.
9. The method of claim 8, wherein the read command further comprises a logical block address of a first logical block of the two or more contiguous logical blocks of data and an entry containing a physical block address of the first logical block.
10. The method of claim 8, further comprising: receiving the set of entries comprising the respective physical block addresses of the two or more contiguous logical blocks of data; and reading the two or more contiguous logical blocks from the respective physical block addresses based at least in part on the transfer length field and the set of entries.
11. The method of claim 8, wherein a size of each of the two or more contiguous logical blocks of data is 4 kilobytes (kB).

12. The method of claim 8, wherein a value of the transfer length field indicates a combined size of the two or more contiguous logical blocks of data.
 13. The method of claim 8, wherein each entry of the set of entries is maintained in a cache of a host device.
 14. The method of claim 8, wherein each entry of the set of entries has a length of 8 bytes.
 15. A memory device, comprising: a memory array; and a controller coupled with the memory array and operable to cause the memory device to: receive a read command comprising: a transfer length field indicating a quantity of two or more contiguous logical blocks of data to be returned; and a value associated with a set of entries that comprise respective physical block addresses of the two or more contiguous logical blocks of data; store the set of entries in association with the value; and transmit the two or more contiguous logical blocks of data read from the respective physical block addresses based at least in part on the transfer length field and the set of entries.
 16. The memory device of claim 15, wherein the read command further comprises a logical block address of a first logical block of the two or more contiguous logical blocks of data and an entry containing a physical block address of the first logical block.
 17. The memory device of claim 15, wherein the controller is further operable to cause the memory device to: receive the set of entries comprising the respective physical block addresses of the two or more contiguous logical blocks of data; and read the two or more contiguous logical blocks from the respective physical block addresses based at least in part on the transfer length field and the set of entries.
 18. The memory device of claim 15, wherein a size of each of the two or more contiguous logical blocks of data is 4 kilobytes (kB).
 19. The memory device of claim 15, wherein a value of the transfer length field indicates a combined size of the two or more contiguous logical blocks of data.
 20. The memory device of claim 15, wherein each entry of the set of entries is maintained in a cache of a host device.
-