

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411798
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Wang; Zixiong

Methods and apparatuses for remote direct memory access page fault handling

Abstract

A method for performing a remote direct memory access (RDMA) read command includes reserving a memory space in a buffer to receive data associated with an address range, sending a RDMA read instruction to a remote module to request the data, and receiving the data from the remote module. A method for performing a RDMA write command includes receiving from a remote module a request to send command requesting to send data associated with an address range, reserving a memory space in a buffer to receive data associated with the address range, sending a RDMA read instruction to the remote module, and receiving the data from the remote module.

Inventors:	Wang; Zixiong (Kanata, CA)
Applicant:	HUAWEI TECHNOLOGIES CO., LTD. (Shenzhen, CN)
Family ID:	94211055
Assignee:	HUAWEI TECHNOLOGIES CO., LTD. (Guangdong, CN)
Appl. No.:	18/220962
Filed:	July 12, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250021515 A1	Jan. 16, 2025

Publication Classification

Int. Cl.: G06F15/17 (20060101); G06F11/07 (20060101); G06F15/173 (20060101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (15/17331); G06F (11/073); G06F (11/0793)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2004/0172521	12/2003	Hooker	712/E9.046	G06F 9/3838
2012/0331065	12/2011	Aho	709/206	G06F 15/167
2020/0097313	12/2019	Dong	N/A	G06F 12/1036

Primary Examiner: Yu; Jae U

Attorney, Agent or Firm: Troutman Pepper Locke LLP

Background/Summary

FIELD OF THE DISCLOSURE

(1) The present disclosure relates generally to virtual computing environments, and in particular to methods and apparatuses for page fault handling for read and write operations using remote direct memory access.

BACKGROUND

(2) Low latency of data transmission has become a key requirement in many critically important computing applications such as data center applications having a cloud-based architecture. Remote direct access memory (RDMA) technologies are commonly used in communication networks to reduce latency and have demonstrated superior latency performance as compared to other technologies such as the use of the transmission control protocol based schemes. However, adoption of RDMA technologies in cloud-based architectures, including for data center applications have been limited due to performance issues as associated memory pages generally require static pinning, which may be unacceptable in many cases due to the resulting effect of poor resource utilization in dynamic multi-tenant environments. Alternatively, the use of on demand paging technologies in cloud-based architectures, including for data center applications may unacceptably increase application latency and degrade application performance.

SUMMARY

(3) Embodiments disclosed herein relate to the use of RDMA read commands to retrieve data from remote computing modules, wherein buffer space in the requesting computing module is reserved in advance of sending a RDMA read command to prevent buffer overflows. Further, embodiments disclosed herein relate to the integration of a RDMA read command in a RDMA write command to similarly prevent buffer overflows in RDMA write operations.

(4) According to one aspect of this disclosure, a method comprises reserving a memory space in a buffer to receive data associated with an address range; sending a RDMA read instruction to a remote module to request the data; and receiving the data from the remote module.

(5) In some embodiments, the method further comprising dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the

RDMA read instruction to request data associated with each address range segment.

(6) In some embodiments, the method further comprises checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

(7) In some embodiments, the method further comprises fetching a work queue element, the work queue element for generating the RDMA read instruction.

(8) In some embodiments, the method further comprises storing a completion event in a completion queue.

(9) According to one aspect of this disclosure, a method comprises receiving from a remote module a request to send command requesting to send data associated with an address range; reserving a memory space in a buffer to receive data associated with the address range; sending a RDMA read instruction to the remote module; and receiving the data from the remote module.

(10) In some embodiments, the method further comprises extracting a tag of the request to send command; checking for a matching work queue element matching the tag; and generating the RDMA read instruction from the matching work queue element.

(11) In some embodiments, wherein the matching work queue element is not found, the method further comprises storing the request to send command in a miss buffer; signaling a container regarding the miss buffer; waiting for confirmation from the container for resolution of the request to send command in the miss buffer; and checking for the matching work queue element matching the tag.

(12) In some embodiments, wherein the request to send command comprises an address offset, a data length, a tag, a source queue pair identification, and a destination queue pair identification.

(13) In some embodiments, the method further comprises extracting the address offset, the data length, the tag, the source queue pair identification, and the destination queue pair identification from the request to send command; checking a tag context table using the tag to obtain an identification of a memory region and a base address of the memory region; and calculating a data start memory address, wherein the data start memory address is the sum of the base address and the address offset.

(14) In some embodiments, the method further comprises sending a completion signal to the remote module.

(15) In some embodiments, the method further comprises checking a page fault bitmap for a page fault, a page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

(16) According to one aspect of this disclosure, an apparatus comprises a memory; and one or more processors for executing instructions stored in the memory for performing RDMA instructions, the instructions comprising reserving memory space in a buffer to receive data associated with an address range; sending a remote direct memory access (RDMA) read instruction to a remote module to request the data; and receiving the data from the remote module.

(17) In some embodiments, the apparatus is further for dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the RDMA read instruction to request data associated with each address range segment.

(18) In some embodiments, the apparatus is further for checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

(19) In some embodiments, the apparatus is further for fetching a work queue element, the work queue element for generating the RDMA read instruction.

(20) According to one aspect of this disclosure an apparatus comprises a memory; and one or more processors for executing instructions stored in the memory for performing RDMA instructions, the instructions comprising receiving from a remote module a request to send command requesting to send data associated with an address range; reserving memory space in a buffer to receive data

associated with the address range; sending a RDMA read instruction to the remote module; and receiving the data from the remote module.

(21) In some embodiments, the apparatus is further for extracting a tag of the request to send command; checking for a matching work queue element matching the tag; and generating the RDMA read instruction from the matching work queue element.

(22) In some embodiments, wherein the matching work queue element is not found, the apparatus is further for storing the request to send command in a miss buffer; signaling a container regarding the miss buffer; waiting for confirmation from the container for resolution of the request to send command in the miss buffer; and checking for the matching work queue element matching the tag.

(23) In some embodiments, the request to send command comprises an address offset, a data length, a tag, a source queue pair identification, and a destination queue pair identification.

(24) In some embodiments, the apparatus is further for extracting the address offset, the data length, the tag, the source queue pair identification, and the destination queue pair identification from the request to send command; checking a tag context table using the tag to obtain an identification of a memory region and a base address of the memory region; and calculating a data start memory address, wherein the data start memory address is the sum of the base address and the address offset.

(25) In some embodiments, the apparatus is further for checking a page fault bitmap for a page fault, a page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) For a more complete understanding of the disclosure, reference is made to the following description and accompanying drawings, in which:

(2) FIG. 1 is a schematic diagram of a conventional computing module in a RDMA network;

(3) FIG. 2 is a schematic diagram of a hypervisor for page fault handling in a RDMA network;

(4) FIG. 3 is a block diagram illustrating conventional RDMA on demand paging;

(5) FIG. 4 is a schematic diagram showing the architecture of a conventional data center;

(6) FIG. 5 is a schematic diagram showing the architecture of a data center according to some embodiments of this disclosure;

(7) FIG. 6 is a schematic diagram of components of a computing module for providing shared virtual address support according to some embodiments of this disclosure;

(8) FIG. 7 is a schematic diagram of a RDMA network over virtual extensible local area network (VxLAN) according to some embodiments of this disclosure;

(9) FIG. 8 is a schematic diagram illustrating a RDMA over VxLAN packet according to some embodiments of this disclosure;

(10) FIG. 9 is a schematic diagram of a computing module of a RDMA network according to some embodiments of this disclosure;

(11) FIG. 10 is the schematic diagram of FIG. 9 further illustrating steps relating to a RDMA read command according to some embodiments of this disclosure;

(12) FIG. 11 is a schematic diagram illustrating steps of a RDMA write command according to some embodiments of this disclosure;

(13) FIG. 12 is the schematic diagram of FIG. 9 further comprising a miss buffer according to some embodiments of this disclosure;

(14) FIG. 13 is the schematic of FIG. 9 further comprising a tag context table according to some embodiments of this disclosure;

(15) FIG. 14 is block diagram of a method for a RDMA read command according to some

embodiments of this disclosure; and

(16) FIG. 15 is block diagram of a method for a RDMA write command according to some embodiments of this disclosure.

DETAILED DESCRIPTION

(17) Remote direct memory access (RDMA) is an extension of direct memory access (DMA) that permits DMA of memory of one computing device by another computing device without involvement of central processing units of the computing devices. The use of RDMA permits applications to directly exchange data using DMA memory semantics without involvement of a host kernel, thus delivering low latency and high performance.

(18) As those skilled in the art will appreciate, the methods, instructions, and processes disclosed herein may be implemented as one or more software and/or firmware programs having necessary computer-executable code or instructions and stored in one or more non-transitory computer-readable storage devices or media which may be any volatile and/or non-volatile, non-removable or removable storage devices such as random access memory, read-only memory, electrically erasable programmable read-only memory, solid-state memory devices, hard disks, compact discs, digital video discs, flash memory devices, and/or the like. Apparatus may read the computer-executable code from the storage devices and execute the computer-executable code to perform the processes.

(19) Alternatively, the methods, instructions, and processes disclosed herein may be implemented as one or more hardware structures having necessary electrical and/or optical components, circuits, logic gates, integrated circuit chips, and/or the like.

(20) Referring to FIG. 1, a RDMA architecture **100** may comprise an application module **102**, a RDMA verb library **108**, a host kernel **110**, and a RDMA network interface card (NIC) **120**. The application module **102** may comprise one or more memory regions **104**. The RDMA verb library **108** may be implemented using a user space verbs library, such as libibverbs, and the host kernel **110** may comprise a verbs core **112**. The NIC **120** may comprise one or more queues **130** for implementing low latency communication and may include a send and receive queue pair (QP) **132** and completion queue (CQ) **138**. The QP **132** may comprise a send queue (SQ) **134** and a receive queue (RQ) **136**. The SQ **134** may be for sending RDMA commands to a remote receiver over a network **140** and the RQ **136** may be for receiving RDMA commands from a remote sender through the network **140**. In some embodiments of the present disclosure, multiple SQs **134** may share a RQ **136** and may be referred to as a share receive queue (SRQ).

(21) The SQ **134** is for storing one or more work queue elements, where each work queue element may be similar to a DMA descriptor, and may be used to generate one or more RDMA commands, such as read, write, and/or the like.

(22) The CQ **138** may be for buffering RDMA command completion events.

(23) The QP **132** and CQ **138** may be in application memory address space, and may be accessed by the NIC **120** and the kernel **110** as well. The QP **132** and the CQ **138** are shown in the NIC **120** in FIG. 1, but in other embodiments, the QP **132** and the CQ **138** may be in host memory as the NIC may not have sufficient memory capacity for such a configuration.

(24) An application may be allocated a memory region **104** for zero-copy access by the NIC **120**, wherein zero-copy access is where the NIC **120** may access memory directly to reduce the number of copies that may be created or used, such as when memory is accessed by the kernel **110**. The memory region **104** may be protected with local and/or remote memory region security keys. The host kernel **110** may maintain association of QPs using a QPC table **114**. The NIC **120** may use a memory address translation table **122** for translating an application's virtual address (va) to a physical memory address.

(25) For each RDMA command, a call initiator refers to a processor and a NIC that generate the RDMA command, while a target refers to a processor and a NIC that receive a RDMA command and execute associated remote data movement. Further, RDMA commands may be split into at least two types of commands, being one-sided commands and two-sided commands. For one-sided

commands, such as RDMA read and write commands, a processor of a target of the RDMA command is not involved in the data movement resulting from the RDMA command. Conversely, for two-sided commands, such as RDMA send and receive commands, a processor of a target of the RDMA command is involved in the data movement results from the RDMA command.

(26) Dynamic memory management is often disabled for memory regions to avoid potential page faults, as RNICs are required to perform read and write accesses to memory regions having unpredictable access times and page locations. For example, if dynamic memory management were used in a system, a memory page in a memory region may be idle and reallocated to another application. Then, when a RNIC begins to access the memory page, a page fault would occur as the memory page was reallocated when it was idle.

(27) Referring to FIG. 2, in some embodiments of the present disclosure, a hypervisor **200** of a server comprises a processor **202**, dynamic random-access memory (DRAM) **216**, and memory storage **222**. The processor **202** comprises a central processing unit (CPU) core **204**, a level 2 (L2) and a level 3 (L3) cache **206**, a L2 and L3 translation lookaside buffer (TLB) **208**, and a memory management unit (MMU), the MMU **210** comprising a page table walker **212**. The DRAM **216** comprises a plurality of pages **216** and a page table **218**, the page table **218** comprising a plurality of page table entries (PTEs) **220**.

(28) Referring to FIG. 2, the hypervisor **200** may maintain a page table **218** for the DRAM **216**, wherein each PTE **220** comprises a present-bit (P-Bit) for indicating whether or not an associated page is in DRAM **216**, or host memory. A page fault occurs when an application attempts to access memory where the PTE **220** has a P-Bit that is clear indicating the associated page is not in the DRAM **216** and an interrupt is raised. The hypervisor **200** may handle a page fault interrupt to resolve it as follows. At step **230**, the CPU core **204** may access the L2 and L3 TLB **208**. At step **232**, if a TLB **208** miss occurs, a page fault interrupt is triggered, and the page table walker **212** resolves the page fault interrupt resulting in long latency for interrupt handling. At step **234**, the page table walker **212** searches an associated PTE **220** in the page table **218**. If the P-Bit in the PTE **220** is set, the page is in host memory and the page address is inserted into the TLB **208** then the page fault interrupt is cleared. Alternatively, at step **236**, if the P-Bit in the PTE **220** is clear, the page is in the storage **222** and swapped from the storage **222** into the DRAM **216**. At step **238**, the P-Bit is then set, the page fault is resolved, and the page fault interrupt is cleared.

(29) Page fault handling in this manner may result in long periods of latency. During latency periods, when a RNIC **120** stores a received packet into a memory page, a page fault may occur as the page is not in a kernel's page table. The received packet would then need to be buffered. However, when the latency for the kernel or RNIC to handle the page fault persists for such a period of time that RNIC does not have adequate buffer capacity, the buffer may also overflow resulting in a buffer overflow. A buffer overflow would result in a loss of RDMA packets because target memory pages are not readily available.

(30) To prevent such packet loss resulting from buffer overflow, all memory pages in a memory region may be statically pinned during a RDMA initialization phase. As a result, the pinned memory pages may not be used by other applications, including when they are idle and not in use. While this addresses the potential for buffer overflow, it may result in significantly inefficient memory use and may be overly restrictive for dynamic data center environments comprising multiple applications and multiple tenants. As a result, RDMA has generally not been widely adopted for data center environments, and has been deployed in a limited set of applications.

(31) One method to address this issue is a RDMA standard comprising on demand paging page fault handling. Referring to FIG. 3, upon receipt of a packet by a RNIC, at step **300**, the RNIC refers to a page table to determine whether a page fault has occurred. If not, normal processing of the packet occurs at step **302**. If a page fault has occurred, the RNIC discards the received packet at step **304**, initiates page fault handling at step **306**, and returns a selective retransmission negative acknowledgement to a remote sender at step **308**. The remote sender would then retransmit the

packet. The use of on demand paging may result in increases to application latency and may degrade application performance.

(32) FIG. 4 is a schematic diagram showing the architecture of a conventional data center **410**. The data center **410** usually comprises a front end dedicated for computing using computer servers and a back end dedicated for storage using storage servers. The front and back ends comprise one or more front-end racks **412** and one or more back-end racks **414**, respectively. As those skilled in the art understand, a data-center rack such as a front-end rack **412** or a back-end rack **414** is a physical structure (such as a steel framework structure) with necessary electronic components and wiring for housing various computing devices such as servers, networking devices, cables, and/or the like, and provides equipment placement and orchestration in the data center facility.

(33) As shown in FIG. 4, each front-end rack **412** comprises a top-of-rack (TOR) switch **422A** (which is a network switch usually placed at the top of the rack) connected to one or more aggregation switches **424A** (which is a network switch for interconnecting multiple racks **412** and **414** in the data center **410**). One or more computer servers **426** are deployed on the front-end racks **412** and connect to the TOR switch **422A** of the front-end rack **412** via an internal or external network interface card (NIC) **428** for accessing a network **432** (such as the Internet) or devices on other racks via the TOR switch **422A** and the aggregation switch **424A**. Each computer sever **426** comprises various computing resources such as one or more central processing units (CPUs) **442**, one or more memory units **444**, one or more storage units **446** such one or more solid-state drives (SSDs), one or more programmable and/or configurable acceleration processors **448** such as one or more graphic processing units (GPUs), one or more DPUs, one or more tensor processing units (TPUs), one or more field-programmable gate arrays (FPGAs), and/or the like (collectively denoted “xPUs”).

(34) Each back-end rack **414** comprises a TOR switch **422B** connected to an aggregation switch **424B**. One or more storage servers **452** are deployed on the back-end racks **414** and connect to the TOR switch **422B** of the back-end rack **414** via an internal or external NIC **428** for accessing the network **432** or devices on other racks via the TOR switch **422B** and the aggregation switch **424B**. Each storage server **452** generally comprises a controller (not shown) controlling one or more non-volatile storage units **454** such as one or more hard drives, one or more SSDs, and/or the like.

(35) In the following, various embodiments of a data-center architecture are disclosed, which provide disaggregation of computing resources with data-centric heterogeneous computing. More specifically, the data-center architecture disclosed below divides hardware resources into physically separated and circuitally interconnected hardware assemblies so that the hardware resources may be more efficiently shared and accessed.

(36) FIG. 5 is a schematic diagram showing the architecture of a data center **500** according to some embodiments of this disclosure. As shown, the data center **500** comprises one or more racks **502**. Each rack **502** comprises a TOR switch **504** connected to one or more aggregation switches **506** for functionally connecting various components, assemblies, and devices on the rack **502** to a network **508** and to the components, assemblies, and devices on other racks.

(37) In these embodiments, the hardware components for computing and depository (similar to those of the computer servers **426** and storage servers **452** shown in FIG. 4) are organized into physically separated and circuitally interconnected hardware assemblies (also denoted “pools” or “nodes”) according to their types. For example, as shown in FIG. 5, the data center **500** comprises one or more CPU nodes **512**, one or more memory nodes **514**, one or more ACC or xPU nodes **516**, and one or more storage nodes **518** for deployment to the one or more racks **502**. Of course, those skilled in the art will appreciate that, in other embodiments, the data center **500** may further comprise other types of nodes as needed.

(38) Each node generally comprises a plurality of same-type hardware components for computing or storage and a DPU, which is a programmable processor specialized for use in data-center infrastructure and streaming applications for processing networking packets, storage requests,

analytics requests, and/or the like) for managing communication of the hardware components of the node with other nodes, other devices, and/or the network **508**. In some embodiments, each node may also comprise a printed circuit board (PCB) or any suitable circuitry having a plurality suitable connection interfaces for receiving the hardware components of the node and establishing communication between the hardware components and between the hardware components and the DPU thereof.

(39) Many different processor architectures comprise input-output memory management unit (IOMMU), system memory management unit (SMMU), and/or the like support for shared virtual address (SVA) technologies, which allow input and output devices such as smart network interface cards (smartNICs), RNICs, and DPUs for accessing memory pages managed by processors with embedded MMUs. Referring to FIG. 6, a processor **600** comprises an integrated MMU **602** and an integrated IOMMU **604**. The MMU **602** and the IOMMU **604** share access to a common guest page table **608** and a common host page table **610** using a common page table walker **606**, all integrated inside the processor **600**.

(40) Embodiments of the present disclosure may be used and is compatible in a variety of applications including for systems using the peripheral component interconnect express (PCIe)/compute express link (CXL) specification. The PCIe/CXL specification supports identification of an application using bus device function (BDF) and software process identification (PASID). When SVA is used, an IOMMU may map tuples comprising BDF and PASID ({BDF, PASID}) into a particular address protection domain.

(41) In some embodiments of the present disclosure, a smartNIC, a RNIC, or a DPU may use SVA to access host DRAM for memory regions. Specifically, a memory may be located in virtual address space of a normal application or guest virtual address space of a virtualized application in a container or a virtual machine (VM).

(42) A smartNIC, a RNIC, or a DPU may access a memory region in a host DRAM for a particular virtual address of a program. A specific memory region may be accessed through the use of tuples comprising the BDF, the PASID, and the virtual address ({BDF, PASID, va}), where the BDF is configured for the associated RNIC or virtualized RNIC in a smartNIC or a DPU.

(43) A smartNIC, a RNIC, or a DPU may also access a memory region in a host DRAM for a particular guest virtual address (gva) of a container or a VM. A specific memory region may be accessed through the use of tuples comprising the BDF, the PASID, and the guest virtual address ({BDF, PASID, gva}), where the BDF is the configured for the associated RNIC or virtualized RNIC in a smartNIC or a DPU.

(44) However, for RDMA, additional memory protections are required for memory region accesses originating from applications in a local servers, also referred to as local RDMA accesses, as well as for memory region accesses originating from applications in remote services, also referred to as remote RDMA accesses. A local RDMA access may occur when the key of the local RDMA access matches a configured local key, 1-key. A remote RDMA access may occur when the key of the remote RDMA access matches a configured remote key, r-key.

(45) To implement the above described memory protections, a memory region may be assigned two PASIDs: a local PASID, PASID-1, for local RDMA accesses, and a remote PASID, PASID-r, for remote RDMA accesses. As a result, for local RDMA accesses to a memory region, a smartNIC, a RNIC, or a DPU may access an associated memory region in host DRAM using a tuples of ({BDF, PASID-1, va}) or ({BDF, PASID-1, gva}). Similarly, for remote RDMA accesses to a memory region, a smartNIC, a RNIC, or a DPU may access an associated memory region in host DRAM using a tuples of ({BDF, PASID-r, va}) or ({BDF, PASID-r, gva}).

(46) In virtualized computing environment such as for a data center or a cloud-based architecture, a tenant may run in multiple containers. To provide high performance and low latency communication for particular applications using such a structure while supporting isolation among tenants, each container may have its own RDMA instance, called a guest RDMA (gRDMA). The

use of gRDMA may be referred to as RDMA virtualization.

(47) Referring to FIG. 7, an RDMA network **700** is configured to operate over an encapsulation protocol, such as virtual extensible local area network (VxLAN). The RDMA network **700** comprises computing resources such as a first computing module **702** and a second computing module **703**, each of the computing modules **702** and **703** comprising containers, hypervisors, and RNICs. The first computing module **702** comprises a first container **704**, a second container **712**, and a third container **714**. The second computing module **703** comprises a fourth container **716**, a fifth container **718**, and a sixth container **720**. The first container **704**, the third container **714**, and the fifth container **718** are for tenant A while the second container **712**, the fourth container **716**, and the sixth container **720** are for tenant B. The first container **702** further comprises a hypervisor **730** and a RNIC **740** and the second container **703** further comprises a hypervisor **726** and a RNIC **728**. Each container **704**, **712**, **714**, **716**, **718**, and **720** is for a designated application with associated memory region and comprises a gRDMA, wherein gRDMA is for communicating with other gRDMA of a common tenant. A group of gRDMA may form a common virtual network instances. For example, the first container **704** is for an application **706** with an associated memory region **708** and comprises a gRDMA **710**. The gRDMA **710** is for communicating with gRDMA of the third container **714** and the fifth container **718**.

(48) Communication between gRDMA may be achieved using RDMA over VxLAN technology, wherein the frame format may be as shown in FIG. 8 in some embodiments disclosed herein. Referring to FIG. 8, a RDMA over VxLAN frame **800** comprises outer headers, a VxLAN header **808**, inner headers, an Infiniband (IB) base transport header (IB BTH+) **816**, a payload **818**, and IB invariant cyclic redundancy check (ICRC) **820**, and an Ethernet frame check sequence (FCS) **824**. The outer headers may comprise an outer media access control (MAC) header **802**, an outer Internet protocol (IP) header **804**, an outer user datagram protocol (UDP) header **806**. The inner headers may comprise an inner MAC header **810**, an inner IP header **812**, and an inner UDP header **814**. In some embodiments disclosed herein, for a particular frame, a container's MAC address may be used for the inner MAC header **810** and the container's IP address may be used for the inner IP header **812**. Further a host's MAC address may be used for the outer MAC header **802** and the host's IP address may be used for the outer IP header **804**.

(49) Referring to FIG. 9, in some embodiments of the present disclosure, a RDMA network **700** comprises a first computing module **702**, the first computing module **702** comprising a first container **704** and a second container **712**, a hypervisor **730**, a RNIC **740**, and a DRAM module **760**. The first container **704** is associated with tenant A and is for an application **706** with an associated memory region **708** and comprises a gRDMA **710**. The hypervisor **730** comprises a verbs core **732**, a MMU **734**, and a IOMMU **436**. The RNIC **740** comprises a response buffer **742** and one or more queues. The RNIC **740** may comprise one or more queues for implementing low latency communication and may include a send and receive QP **744** and CQ **746**. The QP **744** may comprise a SQ **748** and a RQ **750**. In some embodiments of the present disclosure, multiple SQs may share the RQ **750** such that it is referred to as a SRQ. As above, the SQ **748** is for storing one or more work queue elements, where each work queue element may be similar to a DMA descriptor, and may be used to generate one or more RDMA commands, such as read, write, and/or the like. The CQ **746** may be for buffering RDMA command completion events. The DRAM module **760** may comprise a memory region page fault bitmap **762** and a memory region **764**.

(50) The verbs core **732** of the hypervisor **730** and the gRDMA **710** of the first container **704** may receive a page fault interrupt from the memory management unit **734** in the hypervisor **730**. The verbs core **732** and the gRDMA **710** maintain the page fault bitmap **762** in the DRAM module **760** based each of the received page faults. Each bit in the page fault bitmap **762** is associated with a page in the memory region **764**. A set bit in the page fault bitmap **762** means that the associated page has been swapped out of the DRAM module **760**. Alternatively, a clear bit in the page fault bitmap **762** means that the associated page resides in the DRAM module **760**.

- (51) For a RDMA read command initiated by the first computing module **702**, the RNIC **740** reserves space in the response buffer **742**, which may be implemented using static random access memory (SRAM) or high bandwidth memory (HBM). The space in the response buffer **742** is for storing response data resulting from the RDMA read command. The space is reserved prior to sending the command to a target computing module to eliminate the risk of a buffer overflow resulting from the receipt of the data from RDMA read command.
- (52) For a RDMA read command relating to work queue element representing a large amount of data, the RNIC **740** may divide the requested data into multiple segments to reduce the amount of reserved spaced required in the response buffer **742** at any given time. For example, each segment may have a segment size of one kilobyte, five-hundred and twelve bytes, or any other suitable size. The RNIC **740** would then send a RDMA read command serially for each segment.
- (53) The RNIC **740** may optionally check the page fault bitmap for pages overlapping with the memory address range of the work queue element. If the page fault bitmap indicates that a page is swapped out, the RNIC **740** may initiate a read of a first cache line of the page so that the IOMMU may begin resolve potential page faults.
- (54) When a remote RNIC targeted by the RDMA read command receives the read command, it may retrieve the data from the memory region located in the remote RNIC. If a page fault occurs in the remote RNIC, the IOMMU of the processor associated with the remote RNIC may resolve the page fault. Notably, such a page fault in the remote RNIC may result in a delay receiving the data from the RDMA read command due to the page fault but will not result in a buffer overflow.
- (55) Upon receipt of the data requested by the RDMA read command, the requesting RNIC **740** may store the requested data into space reserved in the response buffer **742**. When response data for a RDMA Read command returns to the local RNIC, the local RNIC stores the response data into its reserved space in the response buffer. Then, the RNIC moves the response data from the response buffer into its memory region via IOMMU. Here, IOMMU can resolve any potential page faults, and it won't cause buffer overflow.
- (56) Referring to FIG. **10**, an exemplary embodiment of a RDMA read command is illustrated according to some embodiments disclosed herein. At step **1000**, the requesting RNIC **740** retrieves a work queue element from the send queue **478** and stores the work queue element in the response buffer **742**. At step **1002**, the requesting RNIC **740** reads the page fault bitmap **762** to check for pages overlapping with the address range of the work queue element. At step **1004**, if a page in the address range of the work queue element in the page fault bitmap **762** is set meaning that the page has been swapped out, the RNIC **740** reads a first cache line of that page using the IOMMU. The RNIC **740** may use a tuples in the format of ({BDF, PASIDI, gva}) to access the first cache line. The RNIC **740** then waits for the IOMMU **736** to resolve the resulting page fault error and discards the returned cache line data from the DRAM module **760** as a result of the access of the first cache line. At step **1006**, the requesting RNIC **740** may generate one or more RDMA read commands from the work queue element, reserve space in the response buffer **742** to store received data from the RDMA read commands, then send the RDMA read commands directed towards other targeted computing devices. At step **1008**, the data requested by the RDMA read commands are received from a RNIC of a remote, targeted computing module and are stored in the response buffer **742**. Here, the RNIC of the remote, targeted computing module retrieves the data from a target memory range, resolves any pages faults using its IOMMU, and sends the retrieved data to the requesting RNIC **740**. At step **1010**, the requesting RNIC **740** stores response data into the memory region **764** using the IOMMU **736**, which may address any page fault errors. The address format of tuples ({BDF, PASIDr, gva}) may be used. At step **1012**, the requesting RNIC **740** may store a completion event in the completion queue **746** upon completion of all RDMA read commands from a work queue element.
- (57) A target computing module of a RDMA write command is unable to determine the amount of data associated with the RDMA write command until the command is received. As a result, the

response buffer reservation method described above in relation to RDMA read commands for handling page faults is not immediately suitable for RDMA write commands. In some embodiments disclosed herein, RDMA write commands may incorporate a RDMA read command for avoiding instances of RDMA packet discarding due to response buffer overflow.

(58) Referring to FIG. 11, communication between a RDMA write command source computing module **1100** and a RDMA write command destination **1102** are illustrated. Instead of a RDMA write work queue element from the source computing module **1100** to the destination computing module **1102**, the RDMA write work queue element is converted to a number of discrete steps. First, at step **1104**, a request to send control command (RTS) is sent from the source computing module **1100** to the destination computing module **1102**. In response, at step **1106**, one or more RDMA read commands are sent from the destination computing module **1102** to the source computing module **1100**, and, at step **1108**, read data from one or more RDMA read commands are sent from the source computing module **1100** to the destination computing module **1102**. Then, at step **1110**, a FIN control command is sent from the destination computing module **1102** to the source computing module **1100** indicating that the request to send command has been successfully completed.

(59) In some embodiments disclosed herein, the request to send control command may be implemented as a send command, such a two-sided RDMA command comprising fields such as starting memory address of data source, data length, tag, source and destination queue pair identifiers, or queue pair IDs. Note that the identifier of a memory region, or a memory region ID, and data starting memory address of the destination are not known at the by the data source. However, the processor or container at the data source may have the tag field, and the tag field may therefore be used at the data destination to identify the identifier of the memory region and the data starting memory address at data destination using a process known as tag matching.

(60) RDMA read commands may be used to implement data transfer from the data source to the data destination in this manner. The FIN control command may be implemented as a send command, such as a two-sided RDMA command, to signal completion of the request to send command.

(61) With the integration of a RDMA read command into the RDMA write command, any page faults may be handled through the same mechanisms of the RDMA read command described above. Notably, an important feature of this approach is for the destination computing module **1102** to generate RDMA read commands upon receipt of a request to send command from the source computing module **1100**.

(62) As referenced above, tag matching may be used and is described in more detail as follows. Referring to FIG. 12, first, a hypervisor **730** or container **704** of a computing module at a data destination may post a work queue element to an associated shared receive queue **750**. The work queue element may comprise fields such a memory region identifier, a data start memory address at the memory region, and/or the like. Next, upon receiving a request to send command, the RNIC **740** at a data destination extracts a tag field from the request to send command, and searches the shared receive queue for a work queue element with matching tag field. If a work queue element with a matching tag is located, the RNIC **740** generates one or more RDMA read commands from the fields of the matching work queue element and the request to send command. If a work queue element with a matching tag is not located, the RNIC **740** stores the request to send command into a missing buffer **766** and signals the hypervisor **730** or the container **704**. Once the hypervisor **730** or container **704** posts the associated work queue element to the shared receive queue **750**, the RNIC **740** will extract the tag field from the work queue element and search the miss buffer **766** for a request to send command with a matching tag. If a matching tag is identified, the RNIC **740** generates one or more RDMA read commands from fields of the work queue element and the matching request to send command.

(63) The use of tag matching as described above may be complex and have poor performance at

times as the hypervisor **730** or the container **704** may need to be involved. In some embodiments disclosed herein, tag matching is replaced with the use of offset addresses.

(64) In some embodiments disclosed herein, in addition to fields such as starting memory address at a data source, data length, tag, source and destination queue pair IDs, and/or the like a request to send command comprises an additional field, being an offset address at a destination memory region, which is known to the hypervisor **730** or the container **702** at the data source. Of note, an identifier of memory region of the data destination and the base address of the destination memory region are not known to the hypervisor **730** or the container **702** at the data source.

(65) Referring to FIG. **13** the RNIC **740** further comprises and maintains a tag context table **768**. The tag context table **768** may be indexed either by a tag field or by tuples such as ({destination QP ID, tag}). A entry in the tag context table **768** may store an identifier of an associated memory region at the destination computing module and the base address of the memory region at the destination computing module. The indexing data structure of the tag context table **768** may be implemented in any suitable data structure such as a hash table, radix tree, tree, linear table, and/or the like.

(66) Request to send commands suitable for use with the tag context table **768** may comprise a offset address for the memory region at the destination computing module in addition to other fields, such as start address at the destination computing module, data length, tag, source and destination QP IDs, and/or the like. To generate RDMA read commands from a received request to send command at the destination computing module, the RNIC **740** at the destination computing module may extract the offset address for entry into the tag context table **768** in addition to other fields such as destination QP ID or the tag.

(67) The RNIC **740** may then calculate a starting address in the memory region of the destination computing module being the sum of the memory region base memory and the offset. The RNIC **740** may determine this by referring to the appropriate table entry in the tag context table **768** relating to the request to send command. The RNIC **740** may then generate one or more RDMA read commands using the fields from the request to send command and the calculated starting address in the memory region.

(68) FIG. **14** is a flowchart showing the steps of a method **1400**, according to some embodiments of the present disclosure. The method **1400** begins with, optionally, fetching a work queue element, the work queue element for generating the RDMA read instruction (step **1402**). At step **1404**, optionally, dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the RDMA read instruction to request data associated with each address range segment. At step **1406**, optionally, checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range. At step **1408**, optionally, waiting for resolution of the page fault. At step **1410**, reserving a memory space in a buffer to receive data associated with an address range. At step **1412**, sending a RDMA read instruction to a remote module to request the data. At step **1414**, receiving the data from the remote module. At step **1416**, optionally, resolving page faults in the data. At step **1418**, optionally, storing the data in memory. At step **1420**, optionally, storing a completion event in a completion queue.

(69) FIG. **15** is a flowchart showing steps of a method **1500**, according to some embodiments of the present disclosure. The method **1500** begins with receiving from a remote module a request to send command requesting to send data associated with an address range (step **1502**). At step **1504**, optionally, extracting a tag of the request to send command. At step **1506**, optionally, checking for a matching work queue element matching the tag. At step **1508**, optionally, generating the RDMA read instruction from the matching work queue element. At step **1510**, optionally, storing the request to send command in a miss buffer. At step **1512**, optionally, signaling a container regarding the miss buffer. At step **1514**, optionally, waiting for confirmation from the container for resolution of the request to send command in the miss buffer. At step **1516**, optionally, checking the matching

work queue element matching the tag. At step **1518**, optionally, extracting the address offset, the data length, the tag, the source queue pair identification, and the destination queue pair identification from the request to send command, checking a tag context table using the tag to obtain an identification of a memory region and a base address of the memory region, and calculating a data start memory address. At step **1520**, reserving a memory space in a buffer to receive data associated with the address range. At step **1522**, sending a RDMA read instruction to the remote module. At step **1524**, receiving the data from the remote module. At step **1526**, optionally, storing the data in memory. At step **1528**, optionally, sending a completion signal to the remote module.

(70) Although embodiments have been described above with reference to the accompanying drawings, those of skill in the art will appreciate that variations and modifications may be made without departing from the scope thereof as defined by the appended claims.

Claims

1. A method comprising: receiving from a remote module a request to send command requesting to send data associated with an address range; reserving a memory space in a buffer to receive data associated with the address range; sending a remote direct memory access (RDMA) read instruction to the remote module to request the data; receiving the data from the remote module; extracting a tag of the request to send command; checking for a matching work queue element matching the tag; and generating the RDMA read instruction from the matching work queue element; wherein the matching work queue element is not found, further comprising: storing the request to send command in a miss buffer; signaling a container regarding the miss buffer; waiting for confirmation from the container for resolution of the request to send command in the miss buffer; and checking for the matching work queue element matching the tag.
2. The method of claim 1 further comprising dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the RDMA read instruction to request data associated with each address range segment.
3. The method of claim 1 further comprising: checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.
4. The method of claim 1 further comprising fetching a work queue element, the work queue element for generating the RDMA read instruction.
5. The method of claim 1 further comprising storing a completion event in a completion queue.
6. The method of claim 1 further comprising: extracting an address offset, a data length, a tag, a source queue pair identification, and a destination queue pair identification from the request to send command; checking a tag context table using the tag to obtain an identification of a memory region and a base address of the memory region; and calculating a data start memory address, wherein the data start memory address is the sum of the base address and the address offset, wherein the request to send command comprises the address offset, the data length, the tag, the source queue pair identification, and the destination queue pair identification.
7. The method of claim 1 further comprising sending a completion signal to the remote module.
8. The method of claim 1 further comprising: checking a page fault bitmap for a page fault, a page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.
9. An apparatus comprising: a memory; and one or more processors for executing instructions stored in the memory for performing the method of claim 1.
10. The apparatus of claim 9, wherein the one or more processors are for executing the instructions stored in the memory further for dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the RDMA read

instruction to request data associated with each address range segment.

11. The apparatus of claim 9, wherein the one or more processors are for executing the instructions stored in the memory further for: checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

12. The apparatus of claim 9, wherein the one or more processors are for executing the instructions stored in the memory further for fetching a work queue element, the work queue element for generating the RDMA read instruction.

13. One or more non-transitory computer-readable storage devices comprising computer-executable instructions, wherein the instructions, when executed, cause one or more processors to perform the method of claim 1.

14. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: dividing the address range into two or more address range segments, wherein sending the RDMA read instruction comprises sending the RDMA read instruction to request data associated with each address range segment.

15. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: checking a page fault bitmap for a page fault, the page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.

16. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: fetching a work queue element, the work queue element for generating the RDMA read instruction.

17. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: storing a completion event in a completion queue.

18. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: extracting an address offset, a data length, a tag, a source queue pair identification, and a destination queue pair identification from the request to send command; checking a tag context table using the tag to obtain an identification of a memory region and a base address of the memory region; and calculating a data start memory address, wherein the data start memory address is the sum of the base address and the address offset, wherein the request to send command comprises the address offset, the data length, the tag, the source queue pair identification, and the destination queue pair identification.

19. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: sending a completion signal to the remote module.

20. The one or more non-transitory computer-readable storage devices of claim 13, wherein the instructions, when executed, cause one or more processors to perform: checking a page fault bitmap for a page fault, a page fault where a page in the page fault bitmap overlaps with the address range; and waiting for resolution of the page fault.
