

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411831
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Wang; Xiaobo et al.

Database index performance improvement

Abstract

Managing database operations is provided. The method comprises receiving an insert statement for a database and determining if the insert statement is for a batch insert operation or random insert operation. For a batch insert operation, responsive to determining a threshold number of specified leaf pages are missing from a memory buffer pool of the database, the database asynchronously pre-loads missing leaf pages from a corresponding index on disk into the memory buffer pool. For a random insert operation, responsive to determining a threshold number of specified leaf pages are missing from a memory buffer pool of the database, the database builds at least one memory cache index and inserts key values specified in the insert statement in the memory cache index. The memory cache index is merged with a corresponding index on disk.

Inventors: Wang; Xiaobo (Beijing, CN), Li; Shuo (Beijing, CN), Sun; Sheng Yan (Beijing, CN), Wang; Xiao Hui (Beijing, CN)

Applicant: International Business Machines Corporation (Armonk, NY)

Family ID: 86768106

Assignee: International Business Machines Corporation (Armonk, NY)

Appl. No.: 17/645214

Filed: December 20, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20230195710 A1	Jun. 22, 2023

Publication Classification

Int. Cl.: G06F16/00 (20190101); G06F16/215 (20190101); G06F16/22 (20190101); G06F16/23 (20190101)

U.S. Cl.:

CPC **G06F16/2272** (20190101); **G06F16/215** (20190101); **G06F16/2365** (20190101);

Field of Classification Search

CPC: G06F (16/184); G06F (16/214); G06F (3/061); G06F (3/067); G06F (16/113); G06F (11/14); G06F (3/064); G06F (11/00); G06F (16/211); G06F (16/22); G06F (16/2246); G06F (16/2272); G06F (16/2365); G06F (16/2379); G06F (16/24542); G06F (12/0215); G06F (12/0246); G06F (12/0638); G06F (12/0866); G06F (12/12); G06F (16/9027); G06F (2212/205); G06F (2212/654); G06F (2212/7203); G06F (3/0613); G06F (3/065); G06F (3/0659); G06F (3/0679); G06F (11/2094); G06F (11/2097); G06F (16/2343); G06F (16/27); G06F (16/29); G06F (16/215); G06F (16/2264); G06F (16/2336); G06F (16/2358); G06F (16/2433); G06F (16/9538); G06F (16/9577)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5204958	12/1992	Cheng	707/999.102	G06F 16/9027
10095721	12/2017	Chen et al.	N/A	N/A
10592149	12/2019	Jenkins	N/A	G06F 16/1752
11797510	12/2022	Subramanian Seshadri	N/A	G06F 16/1748
2006/0182030	12/2005	Harris	370/230	H04W 52/04
2015/0347470	12/2014	Ma et al.	N/A	N/A
2016/0092503	12/2015	Deshmukh et al.	N/A	N/A
2019/0205244	12/2018	Smith	N/A	G06F 3/065
2020/0065314	12/2019	Phillips	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
105335482	12/2015	CN	N/A
2023/116347	12/2022	WO	N/A

OTHER PUBLICATIONS

Anonymous, "A Method to Improve Performance for Index Operation," Dec. 31, 2020, An IP.com Prior Art Database Technical Disclosure, IPCOM000264515D, 5 pages. cited by applicant

Anonymous, "Fast and resource efficient table index maintenance following mass table record deletes," Apr. 13, 2011, An IP.com Prior Art Database Technical Disclosure, IPCOM000206068D, 4 pages. cited by applicant

Sadoghi et al., "Making Updates Disk-I/O Friendly Using SSDs," Proceedings of the VLDB Endowment, Aug. 2013, 13 pages. <https://www.researchgate.net/publication/262394013>. cited by applicant

"Prefetching data into the buffer pool," IBM Documentation, IBM Corporation, copyright 2014, accessed Jun. 2, 2021, 4 pages. <https://www.ibm.com/docs/en/db2/11.1?topic=management-prefetching-data-into-buffer-pool>. cited by applicant

Primary Examiner: Hwa; Shyue Jiunn

Attorney, Agent or Firm: Yee & Associates, P.C.

Background/Summary

BACKGROUND

1. Field

(1) The disclosure relates generally to computing systems and more specifically to improving database system performance.

2. Description of the Related Art

(2) Relational databases store and provide access to data in relation to other data. For example, a relational database may allow access to customer data in relation to product orders or vice versa. This data is organized in tables comprising columns and rows which are stored in a number of pages on a disk in the database. In order to delete or update data stored on disk, the pages containing that data must first be loaded into a cache buffer where the operations can be applied to the pages. After the operations are applied to the pages in question, the modified pages are written back to disk.

SUMMARY

(3) An illustrative embodiment provides a computer-implemented method for managing database operations. The method comprises using a number of processors to perform the steps of receiving an insert statement for a database and determining if the insert statement is for a batch insert operation or random insert operation. For a batch insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database, the database asynchronously pre-loads missing leaf pages from a corresponding index on disk in the database into the memory buffer pool. For a random insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database, the database builds at least one memory cache index and inserts key values specified in the insert statement in the memory cache index. The memory cache index is merged with a corresponding index on disk in the database.

(4) According to other illustrative embodiments, a computer system and a computer program product for managing database operations are provided.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 depicts a pictorial representation of a network of data processing systems in which illustrative embodiments can be implemented;

(2) FIG. 2 is a diagram of a data processing system is depicted in accordance with an illustrative embodiment

(3) FIG. 3 depicts a block diagram illustrating a database in accordance with an illustrative embodiment;

(4) FIG. 4 depicts a diagram illustrating a relational database table with which the illustrative embodiments can be implemented;

- (5) FIG. 5 depicts a tree diagram illustrating an index structure with which the illustrative embodiments may be implemented;
- (6) FIG. 6 depicts a diagram illustrating a database buffer pool with which the illustrative embodiments may be implemented;
- (7) FIG. 7 depicts a diagram illustrating database batch Insert in accordance with an illustrative embodiment;
- (8) FIG. 8 depicts a diagram illustrating collecting and sorting of leaf pages for batch Insert in accordance with an illustrative embodiment;
- (9) FIG. 9 depicts a diagram illustrating database random Insert in accordance with an illustrative embodiment;
- (10) FIG. 10 depicts a diagram illustrating consolidating memory cache indexes in accordance with an illustrative embodiment;
- (11) FIG. 11 depicts a diagram illustrating merging a memory cache index into an existing index in accordance with an illustrative embodiment;
- (12) FIG. 12 depicts a diagram illustrating index consolidation range by range in the case of no overlap between indexes in accordance with an illustrative embodiment;
- (13) FIG. 13 depicts a diagram illustrating index consolidation range by range in the case of a small overlap between indexes in accordance with an illustrative embodiment;
- (14) FIG. 14 depicts a diagram illustrating index consolidation page-by-page in the case of a large overlap between indexes in accordance with an illustrative embodiment;
- (15) FIG. 15 depicts a flowchart illustrating a process for managing database operations in accordance with an illustrative embodiment; and
- (16) FIG. 16 depicts a flowchart illustrating a process for merging memory cache indexes in accordance with an illustrating embodiment.

DETAILED DESCRIPTION

- (17) The present invention may be a system, a method, and/or a computer program product at any possible technical detail level of integration. The computer program product may include a computer-readable storage medium (or media) having computer-readable program instructions thereon for causing a processor to carry out aspects of the present invention.
- (18) The computer-readable storage medium can be a tangible device that can retain and store instructions for use by an instruction execution device. The computer-readable storage medium may be, for example, but is not limited to, an electronic storage device, a magnetic storage device, an optical storage device, an electromagnetic storage device, a semiconductor storage device, or any suitable combination of the foregoing. A non-exhaustive list of more specific examples of the computer-readable storage medium includes the following: a portable computer diskette, a hard disk, a random access memory (RAM), a read-only memory (ROM), an erasable programmable read-only memory (EPROM or Flash memory), a static random access memory (SRAM), a portable compact disc read-only memory (CD-ROM), a digital versatile disk (DVD), a memory stick, a floppy disk, a mechanically encoded device such as punch-cards or raised structures in a groove having instructions recorded thereon, and any suitable combination of the foregoing. A computer-readable storage medium, as used herein, is not to be construed as being transitory signals per se, such as radio waves or other freely propagating electromagnetic waves, electromagnetic waves propagating through a waveguide or other transmission media (e.g., light pulses passing through a fiber-optic cable), or electrical signals transmitted through a wire.
- (19) Computer-readable program instructions described herein can be downloaded to respective computing/processing devices from a computer-readable storage medium or to an external computer or external storage device via a network, for example, the Internet, a local area network, a wide area network and/or a wireless network. The network may comprise copper transmission cables, optical transmission fibers, wireless transmission, routers, firewalls, switches, gateway computers and/or edge servers. A network adapter card or network interface in each

computing/processing device receives computer-readable program instructions from the network and forwards the computer-readable program instructions for storage in a computer-readable storage medium within the respective computing/processing device.

(20) Computer-readable program instructions for carrying out operations of the present invention may be assembler instructions, instruction-set-architecture (ISA) instructions, machine instructions, machine dependent instructions, microcode, firmware instructions, state-setting data, configuration data for integrated circuitry, or either source code or object code written in any combination of one or more programming languages, including an object oriented programming language such as Smalltalk, C++, or the like, and procedural programming languages, such as the "C" programming language or similar programming languages. The computer-readable program instructions may execute entirely on the user's computer, partly on the user's computer, as a stand-alone software package, partly on the user's computer and partly on a remote computer, or entirely on the remote computer or server. In the latter scenario, the remote computer may be connected to the user's computer through any type of network, including a local area network (LAN) or a wide area network (WAN), or the connection may be made to an external computer (for example, through the Internet using an Internet Service Provider). In some embodiments, electronic circuitry including, for example, programmable logic circuitry, field-programmable gate arrays (FPGA), or programmable logic arrays (PLA) may execute the computer-readable program instructions by utilizing state information of the computer-readable program instructions to personalize the electronic circuitry, in order to perform aspects of the present invention.

(21) Aspects of the present invention are described herein with reference to flowchart illustrations and/or block diagrams of methods, apparatus (systems), and computer program products according to embodiments of the invention. It will be understood that each block of the flowchart illustrations and/or block diagrams, and combinations of blocks in the flowchart illustrations and/or block diagrams, can be implemented by computer-readable program instructions.

(22) These computer-readable program instructions may be provided to a processor of a computer, or other programmable data processing apparatus to produce a machine, such that the instructions, which execute via the processor of the computer or other programmable data processing apparatus, create means for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks. These computer-readable program instructions may also be stored in a computer-readable storage medium that can direct a computer, a programmable data processing apparatus, and/or other devices to function in a particular manner, such that the computer-readable storage medium having instructions stored therein comprises an article of manufacture including instructions which implement aspects of the function/act specified in the flowchart and/or block diagram block or blocks.

(23) The computer-readable program instructions may also be loaded onto a computer, other programmable data processing apparatus, or other device to cause a series of operational steps to be performed on the computer, other programmable apparatus or other device to produce a computer implemented process, such that the instructions which execute on the computer, other programmable apparatus, or other device implement the functions/acts specified in the flowchart and/or block diagram block or blocks.

(24) The flowcharts and block diagrams in the Figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the present invention. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of instructions, which comprises one or more executable instructions for implementing the specified logical function(s). In some alternative implementations, the functions noted in the blocks may occur out of the order noted in the Figures. For example, two blocks shown in succession may, in fact, be accomplished as one step, executed concurrently, substantially concurrently, in a partially or wholly temporally overlapping manner, or the blocks may sometimes be executed in the reverse order, depending

upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts or carry out combinations of special purpose hardware and computer instructions.

(25) The illustrative embodiments recognize and take into account one or more different considerations. The illustrative embodiments recognize and take into account that in order to delete or update data stored on disk, the pages containing that data must first be loaded into a cache buffer where the operations can be applied to the pages. After the operations are applied to the pages in question, the modified pages are written back to disk.

(26) FIG. 1 depicts a pictorial representation of a network of data processing systems in which illustrative embodiments can be implemented. Network data processing system **100** is a network of computers, data processing systems, and other devices in which the illustrative embodiments may be implemented. Network data processing system **100** contains network **102**, which is the medium used to provide communications links between the computers, data processing systems, and other devices connected together within network data processing system **100**. Network **102** may include connections, such as, for example, wire communication links, wireless communication links, and fiber optic cables.

(27) In the depicted example, server **104** and server **106** connect to network **102**, along with storage **108**. Server **104** and server **106** may be, for example, server computers with high-speed connections to network **102**. In addition, server **104** and server **106** may provide a set of one or more connector services for managing idempotent operations on a system of record, such as storage **108**. An idempotent operation is an identical operation, which was previously performed or executed, that has the same effect as performing a single operation. Also, it should be noted that server **104** and server **106** may each represent a plurality of servers providing management of idempotent operations for a plurality of system of records.

(28) Client **110**, client **112**, and client **114** also connect to network **102**. Clients **110**, **112**, and **114** are clients of server **104** and server **106**. Server **104** and server **106** may provide information, such as boot files, operating system images, and software applications to clients **110**, **112**, and **114**.

(29) In this example, clients **110**, **112**, and **114** are shown as desktop or personal computers. However, it should be noted that clients **110**, **112**, and **114** are intended as examples only. In other words, clients **110**, **112**, and **114** may include other types of data processing systems, such as, for example, network computers, laptop computers, tablet computers, handheld computers, smart phones, smart watches, personal digital assistants, gaming devices, set-top boxes, kiosks, smart tags, IoT sensors, and the like. Users of clients **110**, **112**, and **114** may utilize clients **110**, **112**, and **114** to access system of records corresponding to one or more enterprises, via the connector services provided by server **104** and server **106**, to perform different data operations. The operations may be, for example, retrieve data, update data, delete data, store data, and the like, on the system of records.

(30) Storage **108** is a network storage device capable of storing any type of data in a structured format or an unstructured format. In addition, storage **108** may represent a plurality of network storage devices. Further, storage **108** may represent a system of record, which is an authoritative data source, corresponding to an enterprise, organization, institution, agency, or similar entity. Furthermore, storage **108** may store other types of data, such as authentication or credential data that may include usernames, passwords, and biometric data associated with client users and system administrators, for example.

(31) In addition, it should be noted that network data processing system **100** may include any number of additional servers, clients, storage devices, and other devices not shown. Program code located in network data processing system **100** may be stored on a computer-readable storage medium and downloaded to a computer or other data processing device for use. For example,

program code may be stored on a computer-readable storage medium on server **104** and downloaded to client **110** over network **102** for use on client **110**.

(32) In the depicted example, network data processing system **100** may be implemented as a number of different types of communication networks, such as, for example, an internet, an intranet, a local area network (LAN), and a wide area network (WAN). FIG. **1** is intended as an example only, and not as an architectural limitation for the different illustrative embodiments.

(33) As used herein, when used with reference to items, “a number of” means one or more of the items. For example, “a number of different types of communication networks” is one or more different types of communication networks. Similarly, “a set of,” when used with reference to items, means one or more of the items.

(34) Further, the term “at least one of,” when used with a list of items, means different combinations of one or more of the listed items may be used, and only one of each item in the list may be needed. In other words, “at least one of” means any combination of items and number of items may be used from the list, but not all of the items in the list are required. The item may be a particular object, a thing, or a category.

(35) For example, without limitation, “at least one of item A, item B, or item C” may include item A, item A and item B, or item B. This example may also include item A, item B, and item C or item B and item C. Of course, any combinations of these items may be present. In some illustrative examples, “at least one of” may be, for example, without limitation, two of item A; one of item B; and ten of item C; four of item B and seven of item C; or other suitable combinations.

(36) With reference now to FIG. **2**, a diagram of a data processing system is depicted in accordance with an illustrative embodiment. Data processing system **200** is an example of a computer, such as server **104** in FIG. **1**, in which the illustrative embodiments may be implemented. In this example, data processing system **200** includes communications fabric **202**, which provides communications between processor unit **204**, memory **206**, persistent storage **208**, communications unit **210**, input/output (I/O) unit **212**, and display **214**.

(37) Processor unit **204** serves to execute instructions for software applications and programs that may be loaded into memory **206**. Processor unit **204** may be a set of one or more hardware processor devices or may be a multi-core processor, depending on the particular implementation.

(38) Memory **206** and persistent storage **208** are examples of storage devices **216**. As used herein, a computer-readable storage device or a computer-readable storage medium is any piece of hardware that is capable of storing information, such as, for example, without limitation, data, computer-readable program code in functional form, and/or other suitable information either on a transient basis or a persistent basis. Further, a computer-readable storage device or a computer-readable storage medium excludes a propagation medium, such as transitory signals. Furthermore, a computer-readable storage device or a computer-readable storage medium may represent a set of computer-readable storage devices or a set of computer-readable storage media. Memory **206**, in these examples, may be, for example, a random-access memory (RAM), or any other suitable volatile or non-volatile storage device, such as a flash memory. Persistent storage **208** may take various forms, depending on the particular implementation. For example, persistent storage **208** may contain one or more devices. For example, persistent storage **208** may be a disk drive, a solid-state drive, a rewritable optical disk, a rewritable magnetic tape, or some combination of the above. The media used by persistent storage **208** may be removable. For example, a removable hard drive may be used for persistent storage **208**.

(39) Communications unit **210**, in this example, provides for communication with other computers, data processing systems, and devices via a network, such as network **102** in FIG. **1**.

Communications unit **210** may provide communications through the use of both physical and wireless communications links. The physical communications link may utilize, for example, a wire, cable, universal serial bus, or any other physical technology to establish a physical communications link for data processing system **200**. The wireless communications link may utilize, for example,

shortwave, high frequency, ultrahigh frequency, microwave, wireless fidelity (Wi-Fi), Bluetooth® technology, global system for mobile communications (GSM), code division multiple access (CDMA), second-generation (2G), third-generation (3G), fourth-generation (4G), 4G Long Term Evolution (LTE), LTE Advanced, fifth-generation (5G), or any other wireless communication technology or standard to establish a wireless communications link for data processing system **200**.

(40) Input/output unit **212** allows for the input and output of data with other devices that may be connected to data processing system **200**. For example, input/output unit **212** may provide a connection for user input through a keypad, a keyboard, a mouse, a microphone, and/or some other suitable input device. Display **214** provides a mechanism to display information to a user and may include touch screen capabilities to allow the user to make on-screen selections through user interfaces or input data, for example.

(41) Instructions for the operating system, applications, and/or programs may be located in storage devices **216**, which are in communication with processor unit **204** through communications fabric **202**. In this illustrative example, the instructions are in a functional form on persistent storage **208**. These instructions may be loaded into memory **206** for running by processor unit **204**. The processes of the different embodiments may be performed by processor unit **204** using computer-implemented instructions, which may be located in a memory, such as memory **206**. These program instructions are referred to as program code, computer usable program code, or computer-readable program code that may be read and run by a processor in processor unit **204**. The program instructions, in the different embodiments, may be embodied on different physical computer-readable storage devices, such as memory **206** or persistent storage **208**.

(42) Program code **218** is located in a functional form on computer-readable media **220** that is selectively removable and may be loaded onto or transferred to data processing system **200** for running by processor unit **204**. Program code **218** and computer-readable media **220** form computer program product **222**. In one example, computer-readable media **220** may be computer-readable storage media **224** or computer-readable signal media **226**.

(43) In these illustrative examples, computer-readable storage media **224** is a physical or tangible storage device used to store program code **218** rather than a medium that propagates or transmits program code **218**. Computer-readable storage media **224** may include, for example, an optical or magnetic disc that is inserted or placed into a drive or other device that is part of persistent storage **208** for transfer onto a storage device, such as a hard drive, that is part of persistent storage **208**. Computer-readable storage media **224** also may take the form of a persistent storage, such as a hard drive, a thumb drive, or a flash memory that is connected to data processing system **200**.

(44) Alternatively, program code **218** may be transferred to data processing system **200** using computer-readable signal media **226**. Computer-readable signal media **226** may be, for example, a propagated data signal containing program code **218**. For example, computer-readable signal media **226** may be an electromagnetic signal, an optical signal, or any other suitable type of signal. These signals may be transmitted over communication links, such as wireless communication links, an optical fiber cable, a coaxial cable, a wire, or any other suitable type of communications link.

(45) Further, as used herein, “computer-readable media **220**” can be singular or plural. For example, program code **218** can be located in computer-readable media **220** in the form of a single storage device or system. In another example, program code **218** can be located in computer-readable media **220** that is distributed in multiple data processing systems. In other words, some instructions in program code **218** can be located in one data processing system while other instructions in program code **218** can be located in one or more other data processing systems. For example, a portion of program code **218** can be located in computer-readable media **220** in a server computer while another portion of program code **218** can be located in computer-readable media **220** located in a set of client computers.

(46) The different components illustrated for data processing system **200** are not meant to provide architectural limitations to the manner in which different embodiments can be implemented. In

some illustrative examples, one or more of the components may be incorporated in or otherwise form a portion of, another component. For example, memory **206**, or portions thereof, may be incorporated in processor unit **204** in some illustrative examples. The different illustrative embodiments can be implemented in a data processing system including components in addition to or in place of those illustrated for data processing system **200**. Other components shown in FIG. 2 can be varied from the illustrative examples shown. The different embodiments can be implemented using any hardware device or system capable of running program code **218**.

(47) In another example, a bus system may be used to implement communications fabric **202** and may be comprised of one or more buses, such as a system bus or an input/output bus. Of course, the bus system may be implemented using any suitable type of architecture that provides for a transfer of data between different components or devices attached to the bus system.

(48) The illustrative embodiments recognize and take into account that for a Structured Query Language (SQL) Insert operation, when a new row is inserted into a base table, all related indexes defined on that base table will correspondingly be modified. However, if the index leaf pages being modified are not already in the buffer pool, system performance will be poor due to synchronous read input/output (I/O) to load those pages into the database buffer pool and then perform the relevant modifications against these leaf pages.

(49) The illustrative embodiments recognize and take into account that simply trying to increase the size of the buffer pool or adjusting buffer parameters to improve the page hit ratio has several drawbacks. For one, it is impossible to make the buffer pool size as large as the disk. Furthermore, for pages referenced less times in the buffer pool, performance improvement is quite limited.

(50) The illustrative embodiments also recognize and take into account that dropping the index before a large Insert operation has the drawback of requiring a long time to rebuild the index. In addition, some SQLs cannot execute until the index is available.

(51) The illustrative embodiments provide a method to improve database system performance on Insert, Delete, and Update operations by building a memory cache index to avoid synchronous in page I/O. Insert operations can be performed as batch Inserts or random Insert. For batch Insert, a pre-loading mechanism pre-loads index leaf pages when the buffer pool hit ratio against leaf pages decreases significantly.

(52) For random Insert, the memory cache index is used if leaf pages where the new key value will insert are missed hit in the buffer pool. Multiple memory cache indexes are dynamically merged into one index before consolidation with the existing corresponding index on disk, thereby reducing I/O and improving query performance. The merge may be performed at the page level or range level.

(53) FIG. 3 depicts a block diagram illustrating a database in accordance with an illustrative embodiment. Database **300** might be an example of storage **108** in FIG. 1 and may be implemented in data processing system **200** in FIG. 2. Database **300** comprises instance **302** and disk **330**.

Instance **302** comprises a set of memory structures that manage database files stored on disk **330**.

(54) Instance **302** comprises memory **304** and processes **320**. Memory **304** includes buffer pool **306**, which stores pages **308** that have been read from disk **330**. Database management system **310** comprises software used to manage, retrieve, and manipulate data **332** in the database **300** and makes it possible for end uses to create, read, update, and delete data. Database management system **310** serves as an interface between the database **300** and end users or application programs.

(55) Data management statements **312** received from a client system, e.g., client **110** in FIG. 1, via listener **356** in input/output (I/O) **354** provide instructions to database **300** regarding transactions performed on the data **332** stored on disk **330**. Examples of data management statements **312** include Insert, Delete, and Update. Data management statements **312** might comprises respective ranges of data entries in files **334** to which the transaction applies.

(56) Disk **330** comprises data **332**, which is organized in a number of files **334**. Files **334** in turn comprise a number of tables **336**. Each table **338** comprises a number of columns **340** and rows

342. The actual data rows **342** are stored on a number of pages **346**. Pages **346** constitute the smallest level of I/O and are grouped into extents **344**, typically in groups of eight pages. Files **334** also comprise indexes **348** that are related to specific tables and comprise one or more keys that allow a database query to efficiently retrieve data **332** from disk **330**. Disk **330** also comprises log files **350** that contain transactions which have not yet been committed to the database and control files **352** that keep track of the database's status and physical structure.

(57) The operations comprising a transaction are not performed directly on data **332** on disk **330**. Rather, pages **308** of data are first read into buffer pool **306** from disk **330** before the operations in data management statements **312** can be applied. Processes **320** comprise reading **322** and writing **324** data to and from disk **330**. Logging **326** records transaction information regarding transaction operations applied to any pages **308** in the buffer pool **306**. Checkpoint **328** writes modified (“dirty”) pages from buffer pool **306** to disk **330** and records transaction information in log files **350**.

(58) Table directory **314** contains metadata for the data **332** in database **300**. Table directory **314** comprises a list of all files **334** in database **300**, the number of records in each file, and the names and types of each data field.

(59) Table directory **314** includes memory cache indexes **316** for pages where new key values will insert that are missed in the buffer pool **306**. Table directory **314** also includes a page pre-loader **318** that pre-loads leaf pages from disk **330** into buffer pool **306** for batch Insert operations.

(60) FIG. 4 depicts a diagram illustrating a relational database table with which the illustrative embodiments can be implemented. Table **400** comprises a number of rows **402** and columns **404**. The columns **404** are called attributes and denote the type of data in the table. The rows **402** are called records or tuples and represent instances of the types of data denoted by the columns **404**.

(61) Each row is identified by a primary key in primary keys **406**. Rows in a table can be linked to rows in other tables by adding the primary keys of the linked rows in a column of the other tables. Such keys in other tables that refer to the primary keys of linked rows are known as foreign keys. In the present example, the primary keys **412** in table **410** are listed as foreign keys **408** in table **400**, thereby establishing relationship **414** between the customer table and the product order table.

(62) In practice, database tables are likely to be vastly larger than the simple example depicted in FIG. 4. In such cases, a range operation might span across many pages within a database.

(63) FIG. 5 depicts a tree diagram illustrating an index structure with which the illustrative embodiments may be implemented. In the present example index structure (or B-Tree) **500** is organized into several levels comprising a root node **502**, intermediate levels including nodes **504**, **506**, **508**, and leaf pages/nodes **510**, **512**, **514**, **516**, **518**, **520**. Index structure **500** is an example of how data in a table (e.g., table **400**) is stored in a database.

(64) Index structure **500** is based on the primary key column (e.g., primary keys **406** in FIG. 4) of a database table. Therefore, the data physically stored on disk in a database is sorted according to the primary key column. The actual data is stored in the leaf pages (aka, leaf nodes) **510-520**. The nodes **504**, **506**, **508** between the root node **502** and leaf pages **510-520** are called intermediate levels. In the present example, only one intermediate level is shown for ease of illustration, but there may be multiple intermediate levels depending on the number of rows in the underlying database table.

(65) The leaf pages **510-520** comprise data rows that contain the actual table data and are sorted according to the primary key of the table. The root node **502** and intermediate nodes, nodes **504**, **506**, **508**, comprise index rows. Each index row contains a key value and a pointer to either an intermediate level node or a data row in a leaf page. The pointers help the database find data quickly. Without an index, a database server would have to read every record in the table to find requested data.

(66) A database such as database **300** in FIG. 3 might also employ secondary indexes, which allow efficient access of records by means other than the primary key. Secondary indexes are indexes that

process a segment type in a sequence other than the one defined by the segment's key.

(67) FIG. 6 depicts a diagram illustrating a database buffer pool with which the illustrative embodiments may be implemented. Buffer pool **600** may be an example of buffer pool **306** in FIG. 3. When the database receives an operation statement, the referenced data might be scattered across a large number of pages in the database. Buffer pool **600** might not contain all the pages referenced in the statement. Referenced pages currently in buffer pool **600** are represented by hatched blocks **602**.

(68) In order to complete the operation on referenced pages not in the buffer pool **600**, those pages must first be read from disk **610** into the buffer pool, represented by shaded blocks **604**. However, if the operation references a large number of pages, attempting to read all the pages into the buffer pool **600** at once (synchronously) might place an excessive I/O burden on the database, thereby negatively impacting database performance.

(69) FIG. 7 depicts a diagram illustrating database batch Insert in accordance with an illustrative embodiment. Process **700** may be implemented in database **300** in FIG. 3. During a batch Insert, the database management system loads multiple rows of data into a database table at once rather than inserting them individually.

(70) For a batch Insert, a pre-loading mechanism, such as page pre-loader **318** in FIG. 3, enables dynamically in response to the hit ratio against leaf pages decreasing below a specified threshold. In the present example, the database receives a batch Insert for insertion of columns from Table 2 **702** into Table 1 **704**. However, enough leaf pages from secondary indexes **708** and **710** are missing from buffer pool **706** to trigger the pre-loading mechanism such as page pre-loader **318**.

(71) FIG. 8 depicts a diagram illustrating collecting and sorting of leaf pages for batch Insert in accordance with an illustrative embodiment. Process **800** may be implemented in database **300** in FIG. 3.

(72) For batch Insert, main task **802** is responsible for the database index Insert operation.

Meanwhile, the pre-loading mechanism enables a number of sub-tasks **804**, **806**, **808** that detect whether modified leaf pages from table **810** are missed in the buffer pool **812**. Sub-tasks **804**, **806**, **808** access indexes **814**, **816**, **818** to located leaf pages missed in the buffer pool **812**.

(73) Under the time window concept, sub-tasks **804**, **806**, **808** sort the missing page number and pre-load them into buffer pool **812** via asynchronous I/O for use by main task **802**, thereby avoiding synchronous I/O page loading. The sub-tasks **804**, **806**, **808** stay a step ahead of main task **802**. After batch **820** is pre-loaded into buffer pool **812**, the sub-tasks **804**, **806**, **808** move on to load the next batch **822**.

(74) FIG. 9 depicts a diagram illustrating database random Insert in accordance with an illustrative embodiment. Process **900** may be implemented in database **300** in FIG. 3. During a random Insert, the database management system loads data into random places into a database table rather than contiguous rows.

(75) If leaf pages where new key values are to be inserted are missing from buffer pool **902**, the database dynamically builds a number of memory cache indexes **904** to save the new key values being inserted, thereby reducing insert contention. The number of memory cache indexes built may be based on the size of the random insert. For example, if the number of inserts row exceeds a specified threshold, the inserts may be split into a number of memory cache indexes. The split into separate memory cache indexes may be based on key values according to, e.g., a hash function, key ranges, or a partition index key rule. Non-leaf pages are accessed in existing index **906** on disk to locate the missing leaf pages. The memory cache indexes **904** are then asynchronously merged with the existing index **906** to avoid synchronous I/O page loading.

(76) FIG. 10 depicts a diagram illustrating consolidating memory cache indexes in accordance with an illustrative embodiment. Process **1000** may be implemented in database **300** in FIG. 3.

(77) To reduce I/O overhead and improve query performance of the database, if multiple memory cache indexes were built for a random Insert, they are dynamically merged into a single

consolidated memory cache index before consolidation with an existing index on disk. In the present example, memory cache index **1002** covering leaf pages **300-500** is dynamically merged with memory cache index **1004** covering leaf pages **560-840** to form consolidated memory cache index **1006**. This merging of memory cache indexes **1002** and **1004** occurs within system memory **1010**, which might be an example of memory **304** in FIG. 3. The use of a consolidated memory cache index **1006** for subsequent merging with the index on disk reduces synchronous I/O compared to merging the separate memory cache indexes **1002**, **1004** individually with the existing index on disk.

(78) FIG. 11 depicts a diagram illustrating merging a memory cache index into an existing index in accordance with an illustrative embodiment. Process **1100** may be implemented in database **300** in FIG. 3.

(79) After consolidating separate memory cache indexes, the consolidated memory cache index **1102** is merged with the existing corresponding index **1108** on disk. Leaf pages **1110** from the existing index **1108** covering a range that includes the Insert operations are prefetched from disk and loaded into the buffer pool **1104** of system memory **1106**. The consolidated memory cache index **1102** is then merged into the existing corresponding index **1108** on either a range level or a page level depending on the degree of overlap, if any, between the previously merged memory cached indexes.

(80) FIG. 12 depicts a diagram illustrating index consolidation range by range in the case of no overlap between indexes in accordance with an illustrative embodiment. Process **1200** may be implemented in database **300** in FIG. 3.

(81) In the present example, there is no overlap between the respective leaf pages **1206**, **1212** of the two memory cache indexes **1202**, **1208**. Merging the memory caches indexes **1202**, **1208** into a consolidated memory cache index takes up less room in the buffer pool. In practice, either memory cache index can be merged into the other. A default rule may (but does not have to) specify that a smaller memory cache index is merged into a larger one. Other merge rules may be employed by the illustrative embodiments.

(82) In the present example, the memory cache indexes **1202**, **1208** are merged by copying key values **1214** from the root node **1210** of the second memory cache index **1208** to the root node **1204** of the first memory cache index **1202**. After the key values **1214** are copied from root node **1210** to root node **1204**, the leaf pages **1212** from second memory cache index **1208** are linked to root node **1204**. Key values **1214** are then removed (drained) from root index **1210** before the switch wherein the consolidated memory cache index is merged with the existing local index on disk.

(83) FIG. 13 depicts a diagram illustrating index consolidation range by range in the case of a small overlap between indexes in accordance with an illustrative embodiment. Process **1300** may be implemented in database **300** in FIG. 3.

(84) In a case where there is a small overlap **1306** of leaf pages between the memory cache indexes **1302**, **1304** below a specified threshold, the overlapping leaf pages are first redistributed from one index to the other before consolidating the root nodes. Similar to process **1200** in FIG. 12, the overlapping leaf pages may be redistributed to either memory cache index. Again, a simple default rule may specify that the overlapping pages be redistributed to the larger memory cache index. In the present example, the leaf pages **1314** through **25** in the second memory cache index **1304** are redistributed to the first memory cache index **1302** and then removed from the second memory cache index **1304**, which is also accounted for in the index rows of the intermediate level non-leaf pages **1312**.

(85) After redistributing the overlapping pages, the key values **1316** of the root node **1310** of the second memory cache index **1304** are then copied to the root node **1308** of the first memory cache index **1302** and the intermediate non-leaf pages **1312** are linked to root node **1308**, similar to process **1200** in FIG. 12. The key values **1316** are then removed from root node **1310** before the

consolidated memory cache index is merged with the existing local index on disk.

(86) FIG. 14 depicts a diagram illustrating index consolidation page-by-page in the case of a large overlap between indexes in accordance with an illustrative embodiment. Process 1400 may be implemented in database 300 in FIG. 3.

(87) In the case where there is a large overlap of leaf pages between the memory cache indexes above a specified threshold, the memory cache indexes are consolidated page-by-page. In the present example, the majority of the leaf pages 1410 in the first memory cache index 1402 overlap with the majority of leaf pages 1412 in the second memory cache index 1404. With such a large overlap, trying to simply redistribute the overlap from one memory cache index to another (as in process 1300 in FIG. 13) would be inefficient. Therefore, process 1400 occurs at a lower level of granularity by merging at the page level.

(88) Process 1400 first scans key values in the respective root nodes 1406, 1408 of the first and second memory caches indexes 1402, 1404 and merges them into a new root node 1418. The respective leaf pages 1410, 1412 of each memory cache index are then scanned page-by-page according to the scanned page numbers referenced in the root nodes 1406, 1408. The scanned leaf pages are then sorted into merged leaf pages 1414 thereby eliminating redundant page numbers between leaf pages 1410 and 1412. Corresponding intermediate nodes 1416 may be created as necessary between the new root node 1418 and merged leaf pages 1414. Process 1400 scans the next key values in root nodes 1406, 1408 and continues merging the root nodes into new root node 1418 and scanning leaf pages 1410, 1412 page-by-page to merge them. After root nodes 1406, 1408 have been completely merged into new root node 1418 and leaf pages 1410, 1412 have been sorted into merged leaf pages 1414, root nodes 1406 and 1408 are both drained before the switch to merge the consolidated memory cache index with the existing local index on disk.

(89) In the case of Delete operations, a method similar to that for random Inserts as shown in FIGS. 9-14 can be used but with a Delete index entry placed in the memory cache index instead of an Insert value.

(90) An Update operation would comprise performing an index Delete operation as described above followed by an index random Insert operation.

(91) For a Select statement, the database may use the memory cache index to search first. If the key being search is marked as delete in the memory cache index, nothing is returned, and the index search is stopped. Otherwise, the key value found first is returned, and the Record ID of the key is remembered. The corresponding existing index on disk is then searched. The Record ID from the memory cached index is used for comparison when a key is found in the existing index on disk to avoid duplicate search results.

(92) FIG. 15 depicts a flowchart illustrating a process for managing database operations in accordance with an illustrative embodiment. Process 1500 may be implemented in database 300 shown in FIG. 3 using a data processing system such as data processing system 200 in FIG. 2.

(93) Process 1500 begins by the database receiving an insert statement for a database (step 1502). Process 1500 determines if the insert statement is for a batch insert operation or random insert operation (step 1504). For a batch insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database (step 1506), process 1500 locates the missing leaf pages on disk (step 1508) and asynchronously pre-loads the missing leaf pages from a corresponding index on disk in the database into the memory buffer pool (step 1510). A number of sub-tasks accesses a number of indexes on disk in the database to locate the missing leaf pages and sort them for pre-loading into the memory buffer pool.

(94) For a random insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database (step 1512), process 1500 builds at least one memory cache index (step 1514) and inserts the key values specified in the insert statement in the memory cache index (step 1516). If multiple memory cache

indexes are built, they are consolidated into a single consolidated memory cache index (step **1518**). The memory cache index (or consolidated memory cache index) is then merged with a corresponding index on disk in the database (step **1520**).

(95) Process **1500** then ends.

(96) FIG. **16** depicts a flowchart illustrating a process for merging memory cache indexes in accordance with an illustrating embodiment. Process **1600** may be a more detailed example of step **1518** in FIG. **15**.

(97) Process **1600** begins by determining if there is any overlap of leaf pages referenced by the different memory cache indexes (step **1602**). If there is no overlap of leaf pages, the key values of the first root node are copied to the second root node (step **1604**). The leaf pages of the first memory cache index are then linked to the second root node (step **1606**). The key values are removed (drained) from the first node before the consolidated memory cache index is merged with the existing local index on disk (step **1608**).

(98) If there is an overlap of referenced leaf pages between the memory cache indexes, process **1600** determines if the size of the overlap is over a specified threshold (step **1610**). If the overlap size is below the threshold, the overlapping pages are redistributed from the first memory cache index to the second memory cache index (step **1612**). Process **1600** then proceeds to step **1604**.

(99) If the overlap is above the threshold, the memory cache indexes are merged page-by-page. Key values are scanned in the first root node and second root node (step **1614**) and merged into a new third root node (step **1616**). The leaf pages in the first and second memory cache indexes are then scanned page-by-page according to the scanned key values from the root nodes (step **1618**).

(100) The scanned leaf pages of the first and second memory index caches are merged together, eliminating redundancies between them (step **1620**). The merged leaf pages are linked to the new third root node (step **1622**). The key values are then removed from both the first and second root nodes node before the new consolidated memory cache index is merged with the existing local index on disk (step **1624**).

(101) Process **1600** then ends.

(102) The description of the different illustrative embodiments has been presented for purposes of illustration and description and is not intended to be exhaustive or limited to the embodiments in the form disclosed. The different illustrative examples describe components that perform actions or operations. In an illustrative embodiment, a component can be configured to perform the action or operation described. For example, the component can have a configuration or design for a structure that provides the component an ability to perform the action or operation that is described in the illustrative examples as being performed by the component. Further, to the extent that terms “includes”, “including”, “has”, “contains”, and variants thereof are used herein, such terms are intended to be inclusive in a manner similar to the term “comprises” as an open transition word without precluding any additional or other elements.

(103) The descriptions of the various embodiments of the present invention have been presented for purposes of illustration but are not intended to be exhaustive or limited to the embodiments disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art without departing from the scope and spirit of the described embodiments. The terminology used herein was chosen to best explain the principles of the embodiments, the practical application or technical improvement over technologies found in the marketplace, or to enable others of ordinary skill in the art to understand the embodiments disclosed herein.

Claims

1. A computer-implemented method for managing database operations, the method comprising: using a number of processors to perform the steps of: receiving an insert statement for a database; determining the insert statement is for a batch insert operation or random insert operation; for a

batch insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database, asynchronously pre-loading missing leaf pages from a corresponding index on disk in the database into the memory buffer pool; for a random insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database: building at least one memory cache index; inserting key values specified in the insert statement in the memory cache index; and merging the memory cache index with a corresponding index on disk in the database.

2. The method of claim 1, wherein, for a batch insert, a number of sub-tasks access a number of indexes on disk in the database and locate the missing leaf pages for pre-loading into the memory buffer pool.

3. The method of claim 2, wherein the sub-tasks sort the missing leaf pages for pre-loading into the memory buffer pool.

4. The method of claim 1, wherein building at least one memory cache index comprises building multiple memory cache indexes.

5. The method of claim 4, further comprising consolidating the multiple memory cache indexes into a single consolidated memory cache index prior to merging with the corresponding index on disk.

6. The method of claim 5, wherein consolidating the multiple memory cache indexes comprises merging leaf pages by range when overlap between leaf pages in the indexes is below a specified threshold range.

7. The method of claim 5, wherein consolidating the multiple memory cache indexes comprises merging leaf pages individually when overlap between leaf pages in the indexes is above a specified threshold range.

8. A system for managing database operations, the system comprising: a storage device configured to store program instructions; and one or more processors operably connected to the storage device and configured to execute the program instructions to cause the system to: receive an insert statement for a database; determine the insert statement is for a batch insert operation or random insert operation; for a batch insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database, asynchronously pre-load missing leaf pages from a corresponding index on disk in the database into the memory buffer pool; for a random insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database: build at least one memory cache index; insert key values specified in the insert statement in the memory cache index; and merge the memory cache index with a corresponding index on disk in the database.

9. The system of claim 8, wherein, for a batch insert, a number of sub-tasks access a number of indexes on disk in the database and locate the missing leaf pages for pre-loading into the memory buffer pool.

10. The system of claim 9, wherein the sub-tasks sort the missing leaf pages for pre-loading into the memory buffer pool.

11. The system of claim 8, wherein building at least one memory cache index comprises building multiple memory cache indexes.

12. The system of claim 8, further comprising consolidating the multiple memory cache indexes into a single consolidated memory cache index prior to merging with the corresponding index on disk.

13. The system of claim 12, wherein consolidating the multiple memory cache indexes comprises merging leaf pages by range when overlap between leaf pages in the indexes is below a specified threshold range.

14. The system of claim 12, wherein consolidating the multiple memory cache indexes comprises

merging leaf pages individually when overlap between leaf pages in the indexes is above a specified threshold range.

15. A computer program product for managing database operations, the computer program product comprising: a computer-readable storage medium having program instructions embodied thereon to perform the steps of: receiving an insert statement for a database; determining the insert statement is for a batch insert operation or random insert operation; for a batch insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database, asynchronously pre-loading missing leaf pages from a corresponding index on disk in the database into the memory buffer pool; for a random insert operation, responsive to determining a threshold number of leaf pages specified in the insert statement are missing from a memory buffer pool of the database: building at least one memory cache index; inserting key values specified in the insert statement in the memory cache index; and merging the memory cache index with a corresponding index on disk in the database.

16. The computer program product of claim 15, wherein, for a batch insert, a number of sub-tasks access a number of indexes on disk in the database and locate the missing leaf pages for pre-loading into the memory buffer pool.

17. The computer program product of claim 16, wherein the sub-tasks sort the missing leaf pages for pre-loading into the memory buffer pool.

18. The computer program product of claim 15, wherein the at least one memory cache index comprises multiple memory cache indexes, and further comprising consolidating the multiple memory cache indexes into a single consolidated memory cache index prior to merging with the corresponding index on disk.

19. The computer program product of claim 18, wherein consolidating the multiple memory cache indexes comprises merging leaf pages by range when overlap between leaf pages in the indexes is below a specified threshold range.

20. The computer program product of claim 18, wherein consolidating the multiple memory cache indexes comprises merging leaf pages individually when overlap between leaf pages in the indexes is above a specified threshold range.
