

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411915
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Doiron; Brock et al.

Storing a program on a blockchain

Abstract

A method comprising: retrieving a header file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network. The header file comprises: program header information, and reference information comprising one or more respective transaction IDs of one or more respective second transactions stored on the blockchain. Each second transaction comprises a respective section of a body of the program file, at least one of the sections comprising runnable program code. The method comprises, from the blockchain, retrieving the respective section of the program file from at least one of the second transactions based on the reference information; and running a program comprising the program header information and the at least one retrieved section.

Inventors:	Doiron; Brock (London, GB), Wright; Craig Steven (London, GB)
Applicant:	nChain Holdings Limited (St. John's, AG)
Family ID:	68315353
Assignee:	nChain Licensing AG (Zug, CH)
Appl. No.:	17/642144
Filed (or PCT Filed):	August 21, 2020
PCT No.:	PCT/IB2020/057858
PCT Pub. No.:	WO2021/048665
PCT Pub. Date:	March 18, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20220358189 A1	Nov. 10, 2022

Foreign Application Priority Data

Publication Classification

Int. Cl.: G06F21/12 (20130101); G06F21/60 (20130101); H04L9/00 (20220101)

U.S. Cl.:

CPC G06F21/12 (20130101); G06F21/602 (20130101); H04L9/50 (20220501);

Field of Classification Search

CPC: G06F (21/12); G06F (21/602); H04L (9/50)

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10698728	12/2019	Nelson	N/A	H04L 9/50
10846416	12/2019	Uhr	N/A	G06F 21/608
11201746	12/2020	Ranganathan	N/A	H04L 9/3247
11515999	12/2021	Nelson	N/A	H04L 9/0637
2018/0096349	12/2017	McDonald	N/A	G06Q 20/20
2019/0116024	12/2018	Wright et al.	N/A	N/A
2019/0149337	12/2018	Savanah et al.	N/A	N/A
2019/0199516	12/2018	Carver et al.	N/A	N/A
2019/0392489	12/2018	Tietzen et al.	N/A	N/A
2020/0195442	12/2019	Wright et al.	N/A	N/A
2020/0266985	12/2019	Covaci et al.	N/A	N/A
2020/0351235	12/2019	Shang	N/A	H04L 12/4641
2021/0233074	12/2020	Kramer et al.	N/A	N/A
2022/0051236	12/2021	Wright	N/A	N/A
2022/0261805	12/2021	Tartan et al.	N/A	N/A
2022/0269810	12/2021	Tartan et al.	N/A	N/A
2023/0185767	12/2022	Hung	713/189	G06F 16/152

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
109074565	12/2017	CN	N/A
110692071	12/2019	CN	N/A
111316595	12/2019	CN	N/A
111316615	12/2019	CN	N/A
110062034	12/2020	CN	N/A

2019220148	12/2018	JP	N/A
2020504930	12/2019	JP	N/A
20190076197	12/2018	KR	N/A
20190076197	12/2018	KR	N/A
2009045257	12/2008	WO	N/A
2017145006	12/2016	WO	N/A
2017195160	12/2016	WO	N/A
2017195160	12/2016	WO	N/A
2019008531	12/2018	WO	N/A
2019116248	12/2018	WO	N/A
2020109907	12/2019	WO	N/A
2020109907	12/2019	WO	N/A
2020109908	12/2019	WO	N/A
2020109908	12/2019	WO	N/A
2020109909	12/2019	WO	N/A
2020109909	12/2019	WO	N/A
2020109910	12/2019	WO	N/A
2020109910	12/2019	WO	N/A
2020109911	12/2019	WO	N/A
2020109911	12/2019	WO	N/A
2020109912	12/2019	WO	N/A
2020109912	12/2019	WO	N/A
2020109913	12/2019	WO	N/A
2020109913	12/2019	WO	N/A
2020110025	12/2019	WO	N/A
2020110025	12/2019	WO	N/A
2020212796	12/2019	WO	N/A
2020212796	12/2019	WO	N/A
2021048663	12/2020	WO	N/A
2021048663	12/2020	WO	N/A
WO-2021119210	12/2020	WO	G06Q 20/023

OTHER PUBLICATIONS

Donald E. Knuth, Structured Programming with go to Statements, Computing Surveys, Dec. 1974, pp. 261-301, vol. 6, No. 4, Stanford University, Stanford, California, US. cited by applicant

Andrew Sward et al., Data Insertion in Bitcoin's Blockchain, Mar. 4, 2018, pp. 2379-5980, URL: <http://ledger.pitt.edu/ojs/ledger/article/view/101/93>. cited by applicant

PCT/IB2020/057858, International Search Report and Written Opinion dated Nov. 6, 2020. cited by applicant

GB1913145.7 Search Report dated Mar. 16, 2020. cited by applicant

International Search Report and Written Opinion issued in International Application No. PCT/IB2020/057796, mailed on Nov. 10, 2020, 53 pages. cited by applicant

International Search Report and Written Opinion issued in International Application No. PCT/IB2020/057858, mailed on Nov. 6, 2020, 16 pages. cited by applicant

Knuth D. E., "Structured Programming with go to Statements," Computing Surveys, Stanford University, Stanford, California, US, Dec. 1974, vol. 6, No. 4, 54 pages. cited by applicant

Search Report for the Application No. GB1913143.2 dated Jun. 1, 2020, 14 pages. cited by applicant

Search Report for the Application No. GB1913145.7 dated Mar. 16, 2020, 4 pages. cited by

applicant

Sward A., et al., "Data Insertion in Bitcoin's Blockchain," Mar. 4, 2018, ISSN 2379-5980, 23 pages Retrieved from Internet: URL: <http://ledger.pitt.edu/ojs/ledger/article/view/101/93>. cited by

applicant

Wood G., "Ethereum: A Secure Decentralised Generalised Transaction Ledger Final Draft-Under Review," Jun. 4, 2014, 32 pages, Retrieved from Internet: URL: <http://gavwood.com/paper.pdf>.

cited by applicant

Gupta A., "Cloud Computing Security Using Blockchain," Journal of Emerging Technologies and Innovative Research (JETIR), Jun. 2019, vol. 6(6), pp. 791-794. cited by applicant

Primary Examiner: Torres-Diaz; Lizbeth

Attorney, Agent or Firm: Messner Reeves LLP

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS

(1) This application is the U.S. National Stage of International Application No.

PCT/IB2020/057858 filed on Aug. 21, 2020, which claims the benefit of United Kingdom Patent Application No. 1913145.7, filed on Sep. 12, 2019, the contents of which are incorporated herein by reference in their entireties.

TECHNICAL FIELD

(2) The present disclosure relates to a file format for storing runnable programs in transactions of a blockchain, and for launchings such programs therefrom.

BACKGROUND

(3) A blockchain refers to a form of distributed data structure, wherein a duplicate copy of the blockchain is maintained at each of a plurality of nodes in a peer-to-peer (P2P) network. The blockchain comprises a chain of blocks of data, wherein each block comprises one or more transactions. Each transaction may point back to a preceding transaction in a sequence which may span one or more blocks. Transactions can be submitted to the network to be included in new blocks. New blocks are created by a process known as "mining", which involves each of a plurality of mining nodes competing to perform "proof-of-work", i.e. solving a cryptographic puzzle based on a pool of the pending transactions waiting to be included in blocks.

(4) Each node in the network can have any one, two or all of three roles: forwarding, mining and storage. Forwarding nodes propagate transactions throughout the nodes of the network. Mining nodes perform the mining of transactions into blocks. Storage nodes each store their own copy of the mined blocks of the blockchain. In order to have a transaction recorded in the blockchain, a party sends the transaction to one of the nodes of the network to be propagated. Mining nodes which receive the transaction may race to mine the transaction into a new block. Each node is configured to respect the same node protocol, which will include one or more conditions for a transaction to be valid. Invalid transactions will not be propagated nor mined into blocks. Assuming the transaction is validated and thereby accepted onto the blockchain, then the transaction (including any user data) will thus remain stored at each of the nodes in the P2P network as an immutable public record.

(5) The miner who successfully solved the proof-of-work puzzle to create the latest block is typically rewarded with a new transaction called a "generation transaction" which generates a new amount of the digital asset. The proof-of work incentivises miners not to cheat the system by including double-spending transactions in their blocks, since it requires a large amount of compute

resource to mine a block, and a block that includes an attempt to double spend is likely not be accepted by other nodes.

(6) In an “output-based” model (sometimes referred to as a UTXO-based model), the data structure of a given transaction comprises one or more inputs and one or more outputs. Any spendable output comprises an element specifying an amount of the digital asset, sometimes referred to as a UTXO (“unspent transaction output”). The output may further comprise a locking script specifying a condition for redeeming the output. Each input comprises a pointer to such an output in a preceding transaction, and may further comprise an unlocking script for unlocking the locking script of the pointed-to output. So consider a pair of transactions, call them a first and a second transaction (or “target” transaction). The first transaction comprises at least one output specifying an amount of the digital asset, and comprising a locking script defining one or more conditions of unlocking the output. The second, target transaction comprises at least one input, comprising a pointer to the output of the first transaction, and an unlocking script for unlocking the output of the first transaction.

(7) In such a model, when the second, target transaction is sent to the P2P network to be propagated and recorded in the blockchain, one of the criteria for validity applied at each node will be that the unlocking script meets all of the one or more conditions defined in the locking script of the first transaction. Another will be that the output of the first transaction has not already been redeemed by another, earlier valid transaction. Any node that finds the target transaction invalid according to any of these conditions will not propagate it nor include it for mining into a block to be recorded in the blockchain.

(8) An alternative type of transaction model is an account-based model. In this case each transaction does not define the amount to be transferred by referring back to the UTXO of a preceding transaction in a sequence of past transactions, but rather by reference to an absolute account balance. The current state of all accounts is stored by the miners separate to the blockchain and is updated constantly.

(9) Conventionally the transactions in the blockchain are used to convey a digital asset, i.e. data acting as a store of value. However, a blockchain can also be exploited in order to layer additional functionality on top of the blockchain. For instance, blockchain protocols may allow for storage of additional payload data in a transaction (i.e. user data, or application content, as opposed to data fundamentally relating to how the transaction operates as a transaction on the blockchain network). For instance such payload data may be included an unspendable output of a transaction in an output-based model, e.g. by means of an OP_RETURN opcode included in the locking script which terminates the locking script when run at a node of the blockchain network and thus makes the output non-functional in terms of transacting with the input of another transaction. Modern blockchains are increasing the maximum data capacity that can be stored within a single transaction, enabling more complex data to be incorporated. For instance this may be used to store an electronic document on the blockchain, or even audio or video data.

(10) It is also known to include a small block of executable code in the payload of a transaction. That is, code other than the locking and unlocking script in an output-based model, or other than the operative part of a smart contract in an account-based model. The executable code instead has some other, secondary or “layer 2” (i.e. application level) function. In other words, from the perspective of the transaction protocol employed by the blockchain network, the executable code is just like any other payload data. A client can download the code by inspecting part of the blockchain as stored on one of the storage nodes of the blockchain network, and accessing the code from the payload of the relevant transaction. The code can then be run locally on the client computer.

SUMMARY

(11) However, while it has been known to include a small block of executable code in the payload of a transaction, the inventors are not aware of anyone having previously included a self-contained

program on the blockchain.

(12) Blockchains provide a decentralised and tamper-resistant alternative to conventional database servers with the potential to revolutionize the storage, access, and distribution of data. However, due to its infancy, it lacks the standardized architecture and network protocols to easily build applications and software without extensive knowledge of the underlying technology. For blockchains to reach the point where they are as fundamental an infrastructure as the internet, it is imperative to develop an easy-to-use protocol for accessing data stored on the blockchain by both users and software.

(13) The following discloses a file format suitable for standardizing a range of blockchain-enabled applications. The disclosure establishes a means of storing executable software (i.e. runnable code) on the blockchain that can be reconstructed locally to be run on individual computers with permissioned access. This introduces a new blockchain-based file format for files stored on the blockchain that can be run locally or on a virtual machine.

(14) According to one aspect disclosed herein, there is provided a method of running a program. The method comprises the following steps performed by computer equipment of a consuming party. Firstly, the method comprises retrieving a header file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network. The header file comprises: program header information, and reference information comprising one or more respective transaction IDs of one or more respective second transactions stored on the blockchain. Each of the second transactions comprises a respective section of a body of the program file, at least one of the sections comprising runnable program code. The method then comprises, from the blockchain, retrieving the respective section of the program file from at least one of the one or more second transactions based on the reference information; and running a program comprising the program header information and the at least one retrieved section.

(15) In embodiments, permissions can be encoded in either the transaction storing the file or the file itself in order to allow for control over who can execute the file.

(16) Embodiments may also integrate access to the blockchain into a high-level programming language to facilitate interaction between local software and the blockchain. There may be provided a request-response protocol for accessing data referenced by transaction hyperlinks within blockchain hypertext documents. For instance, the request-response protocol may be facilitated by a hypertext-enabled wallet to make requests via OP_RETURN flags.

(17) The disclosure also discusses an example system architecture and computer architecture, and how components of the computer hardware and software may interact in implementing the disclosed techniques.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) To assist understanding of embodiments of the present disclosure and to show how such embodiments may be put into effect, reference is made, by way of example only, to the accompanying drawings in which:

(2) FIG. 1 is a schematic block diagram of a system for implementing a blockchain,

(3) FIG. 2 schematically illustrates some examples of transactions which may be recorded in a blockchain,

(4) FIG. 3 is a schematic block diagram of an example set of transactions for storing a program,

(5) FIG. 4 is a schematic block diagram of a client application,

(6) FIG. 5 schematically illustrates a blockchain-based file format in accordance with embodiments herein,

(7) FIG. 6 is a schematic block diagram of an example computer architecture for implementing the

presently disclosed techniques, and

(8) FIGS. 7A-7B show an example process for accessing and reconstructing a runnable program file in accordance with embodiments disclosed herein.

DETAILED DESCRIPTION OF EMBODIMENTS

(9) Example System Overview

(10) FIG. 1 shows an example system **100** for implementing a blockchain **150**. The system **100** comprises a packet-switched network **101**, typically a wide-area internetwork such as the Internet. The packet-switched network **101** comprises a plurality of nodes **104** arranged to form a peer-to-peer (P2P) overlay network **106** within the packet-switched network **101**. Each node **104** comprises computer equipment of a peers, with different ones of the nodes **104** belonging to different peers. Each node **104** comprises processing apparatus comprising one or more processors, e.g. one or more central processing units (CPUs), accelerator processors, application specific processors and/or field programmable gate arrays (FPGAs). Each node also comprises memory, i.e. computer-readable storage in the form of a non-transitory computer-readable medium or media. The memory may comprise one or more memory units employing one or more memory media, e.g. a magnetic medium such as a hard disk; an electronic medium such as a solid-state drive (SSD), flash memory or EEPROM; and/or an optical medium such as an optical disk drive.

(11) The blockchain **150** comprises a chain of blocks of data **151**, wherein a respective copy of the blockchain **150** is maintained at each of a plurality of nodes in the P2P network **160**. Each block **151** in the chain comprises one or more transactions **152**, wherein a transaction in this context refers to a kind of data structure. The nature of the data structure will depend on the type of transaction protocol used as part of a transaction model or scheme. A given blockchain will typically use one particular transaction protocol throughout. In one common type of transaction protocol, the data structure of each transaction **152** comprises at least one input and at least one output. Each output specifies an amount representing a quantity of a digital asset belonging to a user **103** to whom the output is cryptographically locked (requiring a signature of that user in order to be unlocked and thereby redeemed or spent). Each input points back to the output of a preceding transaction **152**, thereby linking the transactions.

(12) At least some of the nodes **104** take on the role of forwarding nodes **104F** which forward and thereby propagate transactions **152**. At least some of the nodes **104** take on the role of miners **104M** which mine blocks **151**. At least some of the nodes **104** take on the role of storage nodes **104S** (sometimes also called “full-copy” nodes), each of which stores a respective copy of the same blockchain **150** in their respective memory. Each miner node **104M** also maintains a pool **154** of transactions **152** waiting to be mined into blocks **151**. A given node **104** may be a forwarding node **104F**, miner **104M**, storage node **104S** or any combination of two or all of these.

(13) In a given present transaction **152j**, the (or each) input comprises a pointer referencing the output of a preceding transaction **152i** in the sequence of transactions, specifying that this output is to be redeemed or “spent” in the present transaction **152j**. In general, the preceding transaction could be any transaction in the pool **154** or any block **151**. The preceding transaction **152i** need not necessarily exist at the time the present transaction **152j** is created or even sent to the network **106**, though the preceding transaction **152i** will need to exist and be validated in order for the present transaction to be valid. Hence “preceding” herein refers to a predecessor in a logical sequence linked by pointers, not necessarily the time of creation or sending in a temporal sequence, and hence it does not necessarily exclude that the transactions **152i**, **152j** be created or sent out-of-order (see discussion below on orphan transactions). The preceding transaction **152i** could equally be called the antecedent or predecessor transaction.

(14) The input of the present transaction **152j** also comprises the signature of the user **103a** to whom the output of the preceding transaction **152i** is locked. In turn, the output of the present transaction **152j** can be cryptographically locked to a new user **103b**. The present transaction **152j** can thus transfer the amount defined in the input of the preceding transaction **152i** to the new user

103b as defined in the output of the present transaction **152j**. In some cases a transaction **152** may have multiple outputs to split the input amount between multiple users (one of whom could be the original user **103a** in order to give change). In some cases a transaction can also have multiple inputs to gather together the amounts from multiple outputs of one or more preceding transactions, and redistribute to one or more outputs of the current transaction.

(15) The above may be referred to as an “output-based” transaction protocol, sometimes also referred to as an unspent transaction output (UTXO) type protocol (where the outputs are referred to as UTXOs). A user's total balance is not defined in any one number stored in the blockchain, and instead the user needs a special “wallet” application **105** to collate the values of all the UTXOs of that user which are scattered throughout many different transactions **152** in the blockchain **151**.

(16) An alternative type of transaction protocol may be referred to as an “account-based” protocol, as part of an account-based transaction model. In the account-based case, each transaction does not define the amount to be transferred by referring back to the UTXO of a preceding transaction in a sequence of past transactions, but rather by reference to an absolute account balance. The current state of all accounts is stored by the miners separate to the blockchain and is updated constantly. In such a system, transactions are ordered using a running transaction tally of the account (also called the “position”). This value is signed by the sender as part of their cryptographic signature and is hashed as part of the transaction reference calculation. In addition, an optional data field may also be signed the transaction. This data field may point back to a previous transaction, for example if the previous transaction ID is included in the data field.

(17) With either type of transaction protocol, when a user **103** wishes to enact a new transaction **152j**, then he/she sends the new transaction from his/her computer terminal **102** to one of the nodes **104** of the P2P network **106** (which nowadays are typically servers or data centres, but could in principle be other user terminals). This node **104** checks whether the transaction is valid according to a node protocol which is applied at each of the nodes **104**. The details of the node protocol will correspond to the type of transaction protocol being used in the blockchain **150** in question, together forming the overall transaction model. The node protocol typically requires the node **104** to check that the cryptographic signature in the new transaction **152j** matches the expected signature, which depends on the previous transaction **152i** in an ordered sequence of transactions **152**. In an output-based case, this may comprise checking that the cryptographic signature of the user included in the input of the new transaction **152j** matches a condition defined in the output of the preceding transaction **152i** which the new transaction spends, wherein this condition typically comprises at least checking that the cryptographic signature in the input of the new transaction **152j** unlocks the output of the previous transaction **152i** to which the input of the new transaction points. In some transaction protocols the condition may be at least partially defined by a custom script included in the input and/or output. Alternatively it could simply be a fixed by the node protocol alone, or it could be due to a combination of these. Either way, if the new transaction **152j** is valid, the current node forwards it to one or more others of the nodes **104** in the P2P network **106**. At least some of these nodes **104** also act as forwarding nodes **104F**, applying the same test according to the same node protocol, and so forward the new transaction **152j** on to one or more further nodes **104**, and so forth. In this way the new transaction is propagated throughout the network of nodes **104**.

(18) In an output-based model, the definition of whether a given output (e.g. UTXO) is spent is whether it has yet been validly redeemed by the input of another, onward transaction **152j** according to the node protocol. Another condition for a transaction to be valid is that the output of the preceding transaction **152i** which it attempts to spend or redeem has not already been spent/redeemed by another valid transaction. Again if not valid, the transaction **152j** will not be propagated or recorded in the blockchain. This guards against double-spending whereby the spender tries to spend the output of the same transaction more than once. An account-based model on the other hand guards against double-spending by maintaining an account balance. Because again there is a defined order of transactions, the account balance has a single defined state at any

one time.

(19) In addition to validation, at least some of the nodes **104M** also race to be the first to create blocks of transactions in a process known as mining, which is underpinned by “proof of work”. At a mining node **104M**, new transactions are added to a pool of valid transactions that have not yet appeared in a block. The miners then race to assemble a new valid block **151** of transactions **152** from the pool of transactions **154** by attempting to solve a cryptographic puzzle. Typically this comprises searching for a “nonce” value such that when the nonce is concatenated with the pool of transactions **154** and hashed, then the output of the hash meets a predetermined condition. E.g. the predetermined condition may be that the output of the hash has a certain predefined number of leading zeros. A property of a hash function is that it has an unpredictable output with respect to its input. Therefore this search can only be performed by brute force, thus consuming a substantive amount of processing resource at each node **104M** that is trying to solve the puzzle.

(20) The first miner node **104M** to solve the puzzle announces this to the network **106**, providing the solution as proof which can then be easily checked by the other nodes **104** in the network (once given the solution to a hash it is straightforward to check that it causes the output of the hash to meet the condition). The pool of transactions **154** for which the winner solved the puzzle then becomes recorded as a new block **151** in the blockchain **150** by at least some of the nodes **104** acting as storage nodes **104S**, based on having checked the winner's announced solution at each such node. A block pointer **155** is also assigned to the new block **151_n** pointing back to the previously created block **151_{n-1}** in the chain. The proof-of-work helps reduce the risk of double spending since it takes a large amount of effort to create a new block **151**, and as any block containing a double spend is likely to be rejected by other nodes **104**, mining nodes **104M** are incentivised not to allow double spends to be included in their blocks. Once created, the block **151** cannot be modified since it is recognized and maintained at each of the storing nodes **104S** in the P2P network **106** according to the same protocol. The block pointer **155** also imposes a sequential order to the blocks **151**. Since the transactions **152** are recorded in the ordered blocks at each storage node **104S** in a P2P network **106**, this therefore provides an immutable public ledger of the transactions.

(21) Note that different miners **104M** racing to solve the puzzle at any given time may be doing so based on different snapshots of the unmined transaction pool **154** at any given time, depending on when they started searching for a solution. Whoever solves their respective puzzle first defines which transactions **152** are included in the next new block **151_n**, and the current pool **154** of unmined transactions is updated. The miners **104M** then continue to race to create a block from the newly defined outstanding pool **154**, and so forth. A protocol also exists for resolving any “fork” that may arise, which is where two miners **104M** solve their puzzle within a very short time of one another such that a conflicting view of the blockchain gets propagated. In short, whichever prong of the fork grows the longest becomes the definitive blockchain **150**.

(22) In most blockchains the winning miner **104M** is automatically rewarded with a special kind of new transaction which creates a new quantity of the digital asset out of nowhere (as opposed to normal transactions which transfer an amount of the digital asset from one user to another). Hence the winning node is said to have “mined” a quantity of the digital asset. This special type of transaction is sometime referred to as a “generation” transaction. It automatically forms part of the new block **151_n**. This reward gives an incentive for the miners **104M** to participate in the proof-of-work race. Often a regular (non-generation) transaction **152** will also specify an additional transaction fee in one of its outputs, to further reward the winning miner **104M** that created the block **151_n** in which that transaction was included.

(23) Due to the computational resource involved in mining, typically at least each of the miner nodes **104M** takes the form of a server comprising one or more physical server units, or even whole a data centre. Each forwarding node **104M** and/or storage node **104S** may also take the form of a server or data centre. However in principle any given node **104** could take the form of a user

terminal or a group of user terminals networked together.

(24) The memory of each node **104** stores software configured to run on the processing apparatus of the node **104** in order to perform its respective role or roles and handle transactions **152** in accordance with the node protocol. It will be understood that any action attributed herein to a node **104** may be performed by the software run on the processing apparatus of the respective computer equipment. The node software may be implemented in one or more applications at the application layer, or a lower layer such as the operating system layer or a protocol layer, or any combination of these. Also, the term “blockchain” as used herein is a generic term that refers to the kind of technology in general, and does not limit to any particular proprietary blockchain, protocol or service.

(25) Also connected to the network **101** is the computer equipment **102** of each of a plurality of parties **103** in the role of consuming users. These act as payers and payees in transactions but do not necessarily participate in mining or propagating transactions on behalf of other parties. They do not necessarily run the mining protocol. Two parties **103** and their respective equipment **102** are shown for illustrative purposes: a first party **103a** and his/her respective computer equipment **102a**, and a second party **103b** and his/her respective computer equipment **102b**. It will be understood that many more such parties **103** and their respective computer equipment **102** may be present and participating in the system, but for convenience they are not illustrated. Each party **103** may be an individual or an organization. Purely by way of illustration the first party **103a** is referred to herein as Alice and the second party **103b** is referred to as Bob, but it will be appreciated that this is not limiting and any reference herein to Alice or Bob may be replaced with “first party” and “second party” respectively.

(26) The computer equipment **102** of each party **103** comprises respective processing apparatus comprising one or more processors, e.g. one or more CPUs, GPUs, other accelerator processors, application specific processors, and/or FPGAs. The computer equipment **102** of each party **103** further comprises memory, i.e. computer-readable storage in the form of a non-transitory computer-readable medium or media. This memory may comprise one or more memory units employing one or more memory media, e.g. a magnetic medium such as hard disk; an electronic medium such as an SSD, flash memory or EEPROM; and/or an optical medium such as an optical disc drive. The memory on the computer equipment **102** of each party **103** stores software comprising a respective instance of at least one client application **105** arranged to run on the processing apparatus. It will be understood that any action attributed herein to a given party **103** may be performed using the software run on the processing apparatus of the respective computer equipment **102**. The computer equipment **102** of each party **103** comprises at least one user terminal, e.g. a desktop or laptop computer, a tablet, a smartphone, or a wearable device such as a smartwatch. The computer equipment **102** of a given party **103** may also comprise one or more other networked resources, such as cloud computing resources accessed via the user terminal.

(27) The client application **105** may be initially provided to the computer equipment **102** of any given party **103** on suitable computer-readable storage medium or media, e.g. downloaded from a server, or provided on a removable storage device such as a removable SSD, flash memory key, removable EEPROM, removable magnetic disk drive, magnetic floppy disk or tape, optical disk such as a CD or DVD ROM, or a removable optical drive, etc.

(28) The client application **105** comprises at least a “wallet” function. This has two main functionalities. One of these is to enable the respective user party **103** to create, sign and send transactions **152** to be propagated throughout the network of nodes **104** and thereby included in the blockchain **150**. The other is to report back to the respective party the amount of the digital asset that he or she currently owns. In an output-based system, this second functionality comprises collating the amounts defined in the outputs of the various **152** transactions scattered throughout the blockchain **150** that belong to the party in question.

(29) Note: whilst the various client functionality may be described as being integrated into a given

client application **105**, this is not necessarily limiting and instead any client functionality described herein may instead be implemented in a suite of two or more distinct applications, e.g. interfacing via an API, or one being a plug-in to the other. More generally the client functionality could be implemented at the application layer or a lower layer such as the operating system, or any combination of these. The following will be described in terms of a client application **105** but it will be appreciated that this is not limiting.

(30) The instance of the client application or software **105** on each computer equipment **102** is operatively coupled to at least one of the forwarding nodes **104F** of the P2P network **106**. This enables the wallet function of the client **105** to send transactions **152** to the network **106**. The client **105** is also able to contact one, some or all of the storage nodes **104** in order to query the blockchain **150** for any transactions of which the respective party **103** is the recipient (or indeed inspect other parties' transactions in the blockchain **150**, since in embodiments the blockchain **150** is a public facility which provides trust in transactions in part through its public visibility). The wallet function on each computer equipment **102** is configured to formulate and send transactions **152** according to a transaction protocol. Each node **104** runs software configured to validate transactions **152** according to a node protocol, and in the case of the forwarding nodes **104F** to forward transactions **152** in order to propagate them throughout the network **106**. The transaction protocol and node protocol correspond to one another, and a given transaction protocol goes with a given node protocol, together implementing a given transaction model. The same transaction protocol is used for all transactions **152** in the blockchain **150** (though the transaction protocol may allow different subtypes of transaction within it). The same node protocol is used by all the nodes **104** in the network **106** (though it many handle different subtypes of transaction differently in accordance with the rules defined for that subtype, and also different nodes may take on different roles and hence implement different corresponding aspects of the protocol).

(31) As mentioned, the blockchain **150** comprises a chain of blocks **151**, wherein each block **151** comprises a set of one or more transactions **152** that have been created by a proof-of-work process as discussed previously. Each block **151** also comprises a block pointer **155** pointing back to the previously created block **151** in the chain so as to define a sequential order to the blocks **151**. The blockchain **150** also comprises a pool of valid transactions **154** waiting to be included in a new block by the proof-of-work process. Each transaction **152** (other than a generation transaction) comprises a pointer back to a previous transaction so as to define an order to sequences of transactions (N.B. sequences of transactions **152** are allowed to branch). The chain of blocks **151** goes all the way back to a genesis block (Gb) **153** which was the first block in the chain. One or more original transactions **152** early on in the chain **150** pointed to the genesis block **153** rather than a preceding transaction.

(32) When a given party **103**, say Alice, wishes to send a new transaction **152j** to be included in the blockchain **150**, then she formulates the new transaction in accordance with the relevant transaction protocol (using the wallet function in her client application **105**). She then sends the transaction **152** from the client application **105** to one of the one or more forwarding nodes **104F** to which she is connected. E.g. this could be the forwarding node **104F** that is nearest or best connected to Alice's computer **102**. When any given node **104** receives a new transaction **152j**, it handles it in accordance with the node protocol and its respective role. This comprises first checking whether the newly received transaction **152j** meets a certain condition for being "valid", examples of which will be discussed in more detail shortly. In some transaction protocols, the condition for validation may be configurable on a per-transaction basis by scripts included in the transactions **152**. Alternatively the condition could simply be a built-in feature of the node protocol, or be defined by a combination of the script and the node protocol.

(33) On condition that the newly received transaction **152j** passes the test for being deemed valid (i.e. on condition that it is "validated"), any storage node **104S** that receives the transaction **152j** will add the new validated transaction **152** to the pool **154** in the copy of the blockchain **150**

maintained at that node **104S**. Further, any forwarding node **104F** that receives the transaction **152j** will propagate the validated transaction **152** onward to one or more other nodes **104** in the P2P network **106**. Since each forwarding node **104F** applies the same protocol, then assuming the transaction **152j** is valid, this means it will soon be propagated throughout the whole P2P network **106**.

(34) Once admitted to the pool **154** in the copy of the blockchain **150** maintained at one or more storage nodes **104**, then miner nodes **104M** will start competing to solve the proof-of-work puzzle on the latest version of the pool **154** including the new transaction **152** (other miners **104M** may still be trying to solve the puzzle based on the old view of the pool **154**, but whoever gets there first will define where the next new block **151** ends and the new pool **154** starts, and eventually someone will solve the puzzle for a part of the pool **154** which includes Alice's transaction **152j**). Once the proof-of-work has been done for the pool **154** including the new transaction **152j**, it immutably becomes part of one of the blocks **151** in the blockchain **150**. Each transaction **152** comprises a pointer back to an earlier transaction, so the order of the transactions is also immutably recorded.

(35) Different nodes **104** may receive different instances of a given transaction first and therefore have conflicting views of which instance is 'valid' before one instance is mined into a block **150**, at which point all nodes **104** agree that the mined instance is the only valid instance. If a node **104** accepts one instance as valid, and then discovers that a second instance has been recorded in the blockchain **150** then that node **104** must accept this and will discard (i.e. treat as invalid) the unmined instance which it had initially accepted.

(36) UTXO-Based Model

(37) FIG. 2 illustrates an example transaction protocol. This is an example of an UTXO-based protocol. A transaction **152** (abbreviated "Tx") is the fundamental data structure of the blockchain **150** (each block **151** comprising one or more transactions **152**). The following will be described by reference to an output-based or "UTXO" based protocol. However, this not limiting to all possible embodiments.

(38) In a UTXO-based model, each transaction ("Tx") **152** comprises a data structure comprising one or more inputs **202**, and one or more outputs **203**. Each output **203** may comprise an unspent transaction output (UTXO), which can be used as the source for the input **202** of another new transaction (if the UTXO has not already been redeemed). The UTXO specifies an amount of a digital asset (a store of value). It may also contain the transaction ID of the transaction from which it came, amongst other information. The transaction data structure may also comprise a header **201**, which may comprise an indicator of the size of the input field(s) **202** and output field(s) **203**. The header **201** may also include an ID of the transaction. In embodiments the transaction ID is the hash of the transaction data (excluding the transaction ID itself) and stored in the header **201** of the raw transaction **152** submitted to the miners **104M**.

(39) Say Alice **103a** wishes to create a transaction **152j** transferring an amount of the digital asset in question to Bob **103b**. In FIG. 2 Alice's new transaction **152j** is labelled "Tx.sub.1". It takes an amount of the digital asset that is locked to Alice in the output **203** of a preceding transaction **152i** in the sequence, and transfers at least some of this to Bob. The preceding transaction **152i** is labelled "Tx.sub.0" in FIG. 2. Tx.sub.0 and Tx.sub.1 are just an arbitrary labels. They do not necessarily mean that Tx.sub.0 is the first transaction in the blockchain **151**, nor that Tx.sub.1 is the immediate next transaction in the pool **154**. Tx.sub.1 could point back to any preceding (i.e. antecedent) transaction that still has an unspent output **203** locked to Alice.

(40) The preceding transaction Tx.sub.0 may already have been validated and included in the blockchain **150** at the time when Alice creates her new transaction Tx.sub.1, or at least by the time she sends it to the network **106**. It may already have been included in one of the blocks **151** at that time, or it may be still waiting in the pool **154** in which case it will soon be included in a new block **151**. Alternatively Tx.sub.0 and Tx could be created and sent to the network **102** together, or

Tx.sub.0 could even be sent after Tx.sub.1 if the node protocol allows for buffering “orphan” transactions. The terms “preceding” and “subsequent” as used herein in the context of the sequence of transactions refer to the order of the transactions in the sequence as defined by the transaction pointers specified in the transactions (which transaction points back to which other transaction, and so forth). They could equally be replaced with “predecessor” and “successor”, or “antecedent” and “descendant”, “parent” and “child”, or such like. It does not necessarily imply an order in which they are created, sent to the network **106**, or arrive at any given node **104**. Nevertheless, a subsequent transaction (the descendent transaction or “child”) which points to a preceding transaction (the antecedent transaction or “parent”) will not be validated until and unless the parent transaction is validated. A child that arrives at a node **104** before its parent is considered an orphan. It may be discarded or buffered for a certain time to wait for the parent, depending on the node protocol and/or miner behaviour.

(41) One of the one or more outputs **203** of the preceding transaction Tx.sub.0 comprises a particular UTXO, labelled here UTXO.sub.0. Each UTXO comprises a value specifying an amount of the digital asset represented by the UTXO, and a locking script which defines a condition which must be met by an unlocking script in the input **202** of a subsequent transaction in order for the subsequent transaction to be validated, and therefore for the UTXO to be successfully redeemed. Typically the locking script locks the amount to a particular party (the beneficiary of the transaction in which it is included). I.e. the locking script defines an unlocking condition, typically comprising a condition that the unlocking script in the input of the subsequent transaction comprises the cryptographic signature of the party to whom the preceding transaction is locked.

(42) The locking script (aka scriptPubKey) is a piece of code written in the domain specific language recognized by the node protocol. A particular example of such a language is called “Script” (capital S). The locking script specifies what information is required to spend a transaction output **203**, for example the requirement of Alice's signature. Unlocking scripts appear in the outputs of transactions. The unlocking script (aka scriptSig) is a piece of code written the domain specific language that provides the information required to satisfy the locking script criteria. For example, it may contain Bob's signature. Unlocking scripts appear in the input **202** of transactions.

(43) So in the example illustrated, UTXO.sub.0 in the output **203** of Tx.sub.0 comprises a locking script [Checksig P.sub.A] which requires a signature Sig P.sub.A of Alice in order for UTXO.sub.0 to be redeemed (strictly, in order for a subsequent transaction attempting to redeem UTXO.sub.0 to be valid). [Checksig P.sub.A] contains the public key P.sub.A from a public-private key pair of Alice. The input **202** of Tx.sub.1 comprises a pointer pointing back to Tx.sub.0 (e.g. by means of its transaction ID, TxID.sub.0, which in embodiments is the hash of the whole transaction Tx.sub.0). The input **202** of Tx.sub.1 comprises an index identifying UTXO.sub.0 within Tx.sub.0, to identify it amongst any other possible outputs of Tx.sub.0. The input **202** of Tx.sub.1 further comprises an unlocking script <Sig P.sub.A> which comprises a cryptographic signature of Alice, created by Alice applying her private key from the key pair to a predefined portion of data (sometimes called the “message” in cryptography). What data (or “message”) needs to be signed by Alice to provide a valid signature may be defined by the locking script, or by the node protocol, or by a combination of these.

(44) When the new transaction Tx.sub.1 arrives at a node **104**, the node applies the node protocol. This comprises running the locking script and unlocking script together to check whether the unlocking script meets the condition defined in the locking script (where this condition may comprise one or more criteria). In embodiments this involves concatenating the two scripts:
<Sig P.sub.A> <P.sub.A>||[Checksig P.sub.A]

where “||” represents a concatenation and “< . . . >” means place the data on the stack, and “[. . .]” is a function comprised by the unlocking script (in this example a stack-based language).

Equivalently the scripts may be run one after the other, with a common stack, rather than concatenating the scripts. Either way, when run together, the scripts use the public key P.sub.A of

Alice, as included in the locking script in the output of Tx.sub.0, to authenticate that the locking script in the input of Tx.sub.1 contains the signature of Alice signing the expected portion of data. The expected portion of data itself (the “message”) also needs to be included in Tx.sub.0 order to perform this authentication. In embodiments the signed data comprises the whole of Tx.sub.0 (so a separate element does to need to be included specifying the signed portion of data in the clear, as it is already inherently present).

(45) The details of authentication by public-private cryptography will be familiar to a person skilled in the art. Basically, if Alice has signed a message by encrypting it with her private key, then given Alice's public key and the message in the clear (the unencrypted message), another entity such as a node **104** is able to authenticate that the encrypted version of the message must have been signed by Alice. Signing typically comprises hashing the message, signing the hash, and tagging this onto the clear version of the message as a signature, thus enabling any holder of the public key to authenticate the signature. Note therefore that any reference herein to signing a particular piece of data or part of a transaction, or such like, can in embodiments mean signing a hash of that piece of data or part of the transaction.

(46) If the unlocking script in Tx.sub.1 meets the one or more conditions specified in the locking script of Tx.sub.0 (so in the example shown, if Alice's signature is provided in Tx.sub.1 and authenticated), then the node **104** deems Tx.sub.1 valid. If it is a mining node **104M**, this means it will add it to the pool of transactions **154** awaiting proof-of-work. If it is a forwarding node **104F**, it will forward the transaction Tx.sub.1 to one or more other nodes **104** in the network **106**, so that it will be propagated throughout the network. Once Tx.sub.1 has been validated and included in the blockchain **150**, this defines UTXO.sub.0 from Tx.sub.0 as spent. Note that Tx.sub.1 can only be valid if it spends an unspent transaction output **203**. If it attempts to spend an output that has already been spent by another transaction **152**, then Tx.sub.1 will be invalid even if all the other conditions are met. Hence the node **104** also needs to check whether the referenced UTXO in the preceding transaction Tx.sub.0 is already spent (has already formed a valid input to another valid transaction). This is one reason why it is important for the blockchain **150** to impose a defined order on the transactions **152**. In practice a given node **104** may maintain a separate database marking which UTXOs **203** in which transactions **152** have been spent, but ultimately what defines whether a UTXO has been spent is whether it has already formed a valid input to another valid transaction in the blockchain **150**.

(47) If the total amount specified in all the outputs **203** of a given transaction **152** is greater than the total amount pointed to by all its inputs **202**, this is another basis for invalidity in most transaction models. Therefore such transactions will not be propagated nor mined into blocks **151**.

(48) Note that in UTXO-based transaction models, a given UTXO needs to be spent as a whole. It cannot “leave behind” a fraction of the amount defined in the UTXO as spent while another fraction is spent. However the amount from the UTXO can be split between multiple outputs of the next transaction. E.g. the amount defined in UTXO.sub.0 in Tx.sub.0 can be split between multiple UTXOs in Tx.sub.1. Hence if Alice does not want to give Bob all of the amount defined in UTXO.sub.0, she can use the remainder to give herself change in a second output of Tx.sub.1, or pay another party.

(49) In practice Alice will also usually need to include a fee for the winning miner, because nowadays the reward of the generation transaction alone is not typically sufficient to motivate mining. If Alice does not include a fee for the miner, Tx.sub.0 will likely be rejected by the miner nodes **104M**, and hence although technically valid, it will still not be propagated and included in the blockchain **150** (the miner protocol does not force miners **104M** to accept transactions **152** if they don't want). In some protocols, the mining fee does not require its own separate output **203** (i.e. does not need a separate UTXO). Instead any difference between the total amount pointed to by the input(s) **202** and the total amount of specified in the output(s) **203** of a given transaction **152** is automatically given to the winning miner **104**. E.g. say a pointer to UTXO.sub.0 is the only input

to Tx.sub.1, and Tx.sub.1 has only one output UTXO.sub.1. If the amount of the digital asset specified in UTXO.sub.0 is greater than the amount specified in UTXO.sub.1, then the difference automatically goes to the winning miner **104M**. Alternatively or additionally however, it is not necessarily excluded that a miner fee could be specified explicitly in its own one of the UTXOs **203** of the transaction **152**.

(50) Alice and Bob's digital assets consist of the unspent UTXOs locked to them in any transactions **152** anywhere in the blockchain **150**. Hence typically, the assets of a given party **103** are scattered throughout the UTXOs of various transactions **152** throughout the blockchain **150**. There is no one number stored anywhere in the blockchain **150** that defines the total balance of a given party **103**. It is the role of the wallet function in the client application **105** to collate together the values of all the various UTXOs which are locked to the respective party and have not yet been spent in another onward transaction. It can do this by querying the copy of the blockchain **150** as stored at any of the storage nodes **104S**, e.g. the storage node **104S** that is closest or best connected to the respective party's computer equipment **102**.

(51) Note that the script code is often represented schematically (i.e. not the exact language). For example, one may write [Checksig P.sub.A] to mean [Checksig P.sub.A]=OP_DUP OP_HASH160 <H(P.sub.A)>OP_EQUALVERIFY OP_CHECKSIG. "OP_ . . ." refers to a particular opcode of the Script language. OP_CHECKSIG (also called "Checksig") is a Script opcode that takes two inputs (signature and public key) and verifies the signature's validity using the Elliptic Curve Digital Signature Algorithm (ECDSA). At runtime, any occurrences of signature ('sig') are removed from the script but additional requirements, such as a hash puzzle, remain in the transaction verified by the 'sig' input. As another example, OP_RETURN is an opcode of the Script language for creating an unspendable output of a transaction that can store metadata within the transaction, and thereby record the metadata immutably in the blockchain **150**. E.g. the metadata could comprise a document which it is desired to store in the blockchain.

(52) The signature P.sub.A is a digital signature. In embodiments this is based on the ECDSA using the elliptic curve secp256k1. A digital signature signs a particular piece of data. In embodiments, for a given transaction the signature will sign part of the transaction input, and all or part of the transaction output. The particular parts of the outputs it signs depends on the SIGHASH flag. The SIGHASH flag is a 4-byte code included at the end of a signature to select which outputs are signed (and thus fixed at the time of signing).

(53) The locking script is sometimes called "scriptPubKey" referring to the fact that it comprises the public key of the party to whom the respective transaction is locked. The unlocking script is sometimes called "scriptSig" referring to the fact that it supplies the corresponding signature. However, more generally it is not essential in all applications of a blockchain **150** that the condition for a UTXO to be redeemed comprises authenticating a signature. More generally the scripting language could be used to define any one or more conditions. Hence the more general terms "locking script" and "unlocking script" may be preferred.

(54) Client Software

(55) FIG. 4 illustrates an example implementation of the client application **105** for implementing embodiments of the presently disclosed techniques. The client application **105** comprises a transaction engine **404** and a user interface (UI) layer **401**. The transaction engine **404** is configured to implement the underlying transaction-related functionality of the client **105**, such as to formulate transactions **152**, receive and/or send transactions and/or other data over the side channel **301**, and/or send transactions to be propagated through the P2P network **106**, in accordance with the schemes discussed above. Furthermore, in accordance with embodiments disclosed herein, the transaction engine **401** of at least Bob's client **105b** (and potentially other clients) further comprises a data transfer engine **403**, a program launcher **405**, and optionally a hyperdocument engine **402**. Together the hyperdocument engine **402**, data transfer engine **403** and transaction engine **404** may form a "smart wallet" function **406** (though the hyperdocument engine **402** is an optional part of

this). The functionality of these various modules will be discussed in more detail shortly. The smart wallet function **406** may also be referred to herein as just the “smart wallet” for short.

(56) The UI layer **401** is configured to render a user interface via a user input/output (I/O) means of the respective user's computer equipment **102**, including outputting information to the respective user **103** via a user output means of the equipment **102**, and receiving inputs back from the respective user **103** via a user input means of the equipment **102**. For example the user output means could comprise one or more display screens (touch or non-touch screen) for providing a visual output, one or more speakers for providing an audio output, and/or one or more haptic output devices for providing a tactile output, etc. The user input means could comprise for example the input sensor matrix of one or more touch screens (which could be the same or different as the display used for the output means); one or more cursor-based devices such as mouse, trackpad or trackball; one or more microphones and speech or voice recognition algorithms for receiving a speech or vocal input; one or more gesture-based input devices for receiving the input in the form of manual or bodily gestures; or one or more mechanical buttons, switches or joysticks, etc.

(57) The transaction engine **403** of the smart wallet **406** is operable to access the code of a stored in the payload of transactions **152** on the blockchain **150** (e.g. in an unspendable output in an output-based transaction model), and passes the code therefrom to the program launcher **405**. The code launcher **405** runs this code and surfaces any user-visible effects of the program to the user **103b** through the UI layer **401**. Particularly, in accordance with embodiments disclosed herein, the code launcher **405** is configured to reconstruct a program file from multiple transactions **152** on the blockchain **150**, and run the reconstructed program file. This will be discussed in more detail shortly.

(58) Note: whilst the various functionality herein may be described as being integrated into the same client application **105**, this is not necessarily limiting and instead they could be implemented in a suite of two or more distinct applications, e.g. one being a plug-in to the other or interfacing via an API (application programming interface). For instance, the functionality of any of the modules **401**, **402**, **403**, **404**, **405** may be implemented in two or more separate applications in any combination. Nor is it excluded that some or all of the described functionality could be implemented at, say, the operating system layer. Where reference is made anywhere herein to a single or given application **105**, or such like, it will be appreciated that this is just by way of example, and more generally the described functionality could be implemented in any form of software.

(59) Blockchain-Based Program File Format

(60) To support the integration of blockchain technology as a fundamental technological infrastructure, there should preferably be a straightforward process to build and access higher-layer applications. In order to do this, it would be desirable to provide standard data formats for developers to not only include within transactions but to build software to execute code stored within blockchain transactions. This may involve dedicated smart wallets that can readily interact with the blockchain both for data requests and to facilitate the connection of software on a local computer, a central processing unit (CPU) and computer storage. Through being able to integrate blockchain data into conventional programming languages, a wider range of applications can be developed atop the blockchain powering more secure and tamper-proof desktop and mobile applications.

(61) The system disclosed herein facilitates direct interaction between software and the blockchain by providing for a transaction-enabled file format for containing programming languages other than the scripting language used for the locking and unlocking scripts or the like. In contemporary programming languages such as C, C++, and Java, there are objects called pointers that act as references to particular memory addresses. The present disclosure provides an analogous object which may be called a blockchain pointer, that can be integrated into programming languages that acts as a reference to some transaction data or metadata.

(62) Some embodiments may also employ a system for the request and response to data requests via linked hypertext documents referenced by their TxID.

(63) FIG. 5 illustrates a blockchain-based program file format for storing a program file across multiple transactions **152** on a blockchain **150** in accordance with embodiments disclosed herein. The left hand panel a) of FIG. 5 shows the program file in fragmented form as stored across multiple transactions **152** on the blockchain **150**. The fragmented form of the program file comprises a header file **501** which is stored in a first transaction **152Y** on the blockchain **150**, and one or more body sections **502_i . . . 502_n** each of which is stored in a respective second transaction **152Z** on the blockchain (see also FIG. 6, to be discussed in more detail later). In an output-based model, each of the header file **501** and body sections **502** may be stored in an un-spendable output (e.g. UTXO) of its respective first or second transaction **152Y**, **152Z**.

(64) An example of this is illustrated schematically in FIG. 3. In the following example scenarios Alice **103a** is the providing party who formulates the first and second transactions **152Y**, **152Z** and broadcasts them to the blockchain network **106** (either directly or via a third party) in order for them to be recorded on the blockchain **150**. Bob **103b** is the consuming party who reads the header file **501** and sections **502** from the respective first and second transactions **152Y**, **152Z** in order to reconstruct the program file **507**.

(65) As shown in FIG. 3, in an output-based model the first transaction **152Y** comprises at least one input **202Y** pointing back to the output of some other, antecedent transaction **152** on the blockchain **150**. The output of the antecedent transaction could be an output locked to Alice, i.e. so Alice is paying an amount of the digital asset from herself to have the first transaction **152Y** recorded on the blockchain **150**. The first transaction **152Y** will typically also be required to comprise at least one spendable output **203Y.sub.0**. This will comprise a locking script locking the output **203Y.sub.0** to a payee. The locking script is formulated in the relevant scripting language recognised by the transaction protocol being employed by the block chain network **106**, e.g. the Script language. The amount specified in this spendable output **203Y.sub.0** could be a dust (i.e. negligible) amount and/or the payee could be Alice herself, i.e. so Alice is paying herself (plus the mining fee) to have the first transaction **152Y** recorded on the blockchain **150**. Alternatively it is not excluded that a transaction model or protocol could be employed that does not necessarily require a spendable output to be included in every transaction. Either way, the first transaction **152Y** also comprises an un-spendable output **203Y.sub.1**. In embodiments, this is implemented by including a special opcode of the scripting language in the locking script of the un-spendable output **203Y.sub.1**. This is an opcode which, when run on one of the nodes **104** of the blockchain network **106**, of terminating execution of the locking script in which it is included. This enables any other arbitrary payload, not formulated in the scripting language, to be included in the un-spendable output **203Y.sub.1** without otherwise affecting the functionality of the transaction **152Y** as a transaction on the blockchain network **106**.

(66) In accordance with embodiments disclosed herein, this mechanism is used to include the header file **501** in the un-spendable output **203Y.sub.1** of the first transaction **152Y**, and thus store the header file **501** in the first transaction **152Y** on the blockchain **150**. The header file **501** is formulated in a format or language other than the scripting language (e.g. Script) used in the locking scripts according to the transaction protocol employed by the blockchain network **106**.

(67) An analogous mechanism may be used to include each of the body sections **502** in an un-spendable output **203Z.sub.1** of a respective second transaction **152Z** on the blockchain **150**. Each body portion **502** comprises program code formulated in a language other than the scripting language (e.g. Script) used in the locking scripts.

(68) In alternative implementations, it is not excluded that the header file **501** and body sections **502** could be included as transaction payloads according to a different mechanism in other transaction protocols or models. For instance, in an alternative implementation, the blockchain network **106** could employ an account-based model, and each of the header file **501** and body

portions **502** could be included in a data field or smart contract of a respective transaction **152Y**, **152Z**. In this case, each of the header file **501** and body sections **502** would be formulated in a language other than that recognised by the nodes **104** for executing smart contracts.

(69) Returning to FIG. 5, the header file **501** comprises at least a portion of program header information **503**, and a list of one or more transaction IDs **504**. The program header information **503** comprises basic header information for the program file that will form at least part of a header of the program file **507** once reconstructed. This may comprise for example any one, more or all of: a version number of the program, a version number of a target computer system, a size of the program or reconstructed program file, a description, reserved variables, a date/time stamp, and/or a checksum for checking the reconstructed program is free of errors.

(70) The transaction IDs included in the list **504** act as pointers to respective body portions **502** of the program file stored in respective other, second transactions **152Z** on the blockchain **150**. As shown in the right hand panel b) of FIG. 5, the consuming party Bob **103b** reconstructs an instance of the program file **507** in un-fragmented form by combining the header information header **503** with each of the one or more body sections **502i . . . 502n** pointed to by the respective transaction IDs **504** included in the header file **501**.

(71) In embodiments, the header file **501** may also comprise an indication of one or more conditions **505**, and a code stub **506** for evaluating the one or more conditions. The functionality of these portions will be discussed in more detail later.

(72) FIG. 6 shows an example computer architecture of the computer equipment **102b** of the consuming party **103b**, also referred to herein as the client equipment. This architecture may be implemented locally on a single user device, or may be distributed across multiple devices which could comprise user devices, server resources or a combination thereof.

(73) The client equipment **102b** comprises hardware **608**, input/output (I/O) devices **607**, and software. The hardware **608** comprises at least one processor **609** such as a CPU. The hardware **608** also comprises random access memory (RAM) and hard storage **611** (i.e. non-volatile memory such as a hard disk or solid state drive, SSD). The central processing unit (CPU) **609** comprises electronic circuitry that carries out machine-level instructions by performing arithmetic logic as well as input/output operations. The RAM **610** provides temporary storage of data and machine-level instructions acting as in intermediary between the software and hard disk **611** with the CPU **609**.

(74) The user **103b** can interact with the input/output devices **607** and the commands are translated to instructions/data, which are relayed to software in the operating system **602**. Various options for I/O devices **607** have already been discussed previously with reference to FIG. 4.

(75) The software comprises an operating system **602**, and a program launcher **405** and smart wallet function **406** that run on the operating system **602**. The operating system **602** comprises system software that manages computer hardware and software resources.

(76) The following definitions may also be adopted herein.

(77) Virtual machine: emulation of a computer system.

(78) System virtual machine: a virtual machine that provides functionality to execute entire operating systems. A hypervisor uses native execution (machine code, directly executed by a computer's CPU) to share and manage hardware, allowing for multiple isolated environments to exist on a single physical machine.

(79) Hypervisor: computer software, firmware, or hardware that creates and runs virtual machines.

(80) Dynamic Linker: part of an operating system that loads and links the shared libraries needed by an executable, when it is executed ("at run time"), by copying the content of libraries from persistent storage to RAM and filling up jump tables and relocating pointers. The specific operating system and executable format determine how the dynamic linker functions and how it is implemented.

(81) Library: a collection of resources used by computer programs (configuration data,

documentation, pre-written code as well as subroutines, classes, values or type specifications).

(82) Shared Library: a file that is shared by executable files and further shared object files.

Modules used are loaded from individual shared objects into memory at run time rather than being copied by a linker when it creates a single monolithic executable file for the program.

(83) Assembler: a computer program that creates object code by translating combinations of mnemonics and syntax for operations and addressing modes into their numerical equivalence (machine code).

(84) Compiler: a computer program that transforms computer code written in one programming language into another (typically from a high-level to lower level language) to create an executable program.

(85) Executable file or executable program: a file that causes a computer to perform indicated tasks according to encoded instructions. Note also that most generally, the term “execute” or “executable” herein can refer to any manner of running software, and does not necessarily limit to execution of low-level machine-code instructions or such like, for example. E.g. “execute” herein could also refer to compiling and then running a high-level language, assembling and then running an assembly level language, or interpreting an interpreted language.

(86) FIG. 6 also illustrates transactions **152** of the blockchain **150** stored on the blockchain network **106**, including the first transaction **152Y** and the second transaction(s) **152Z**.

(87) In operation, the consuming party Bob **103b** uses at least one of the I/O devices **607** to select to access the header file **501** from the first transaction **152Y** as stored on the blockchain **150** (it will be appreciated that “first” and “second” as used herein in this context are just arbitrary labels). In some embodiments, this may be achieved by selecting a hyperlink in a hypertext document **601**. For instance, this may be formulated in HTML, but augmented with an additional type of tag that enables a link to a transaction ID of a transaction **152** on the blockchain **150**, instead of to an URL as in conventional hyperlinks. In such embodiments the hypertext engine **402** of the smart wallet **406** renders the hypertext document **601** to Bob through the UI layer **401**, and detects the selection of the hyperlink. In response, it triggers the data transfer engine **403** of the wallet **406** to retrieve the header file **501** from the linked-to transaction (the first transaction **152Y**) on the blockchain **150** via the blockchain network **106** (by querying a storage node **104S**).

(88) The hypertext document **601** could be accessed from a local store on Bob's computer equipment **102b** or via the Web. Alternatively, the hypertext document **601** itself could be stored as payload data in a source transaction **152X** on the blockchain **150**, e.g. in an unspendable output in an output-based model. In this case, in a preliminary step the data transfer engine **403** accesses the hypertext document **601** from the source transaction **152X** on the blockchain via the blockchain network **106**, and passes it to the hypertext engine **402** to render to the user **103b**.

(89) In alternative embodiments, it is not necessary to use a hypertext document **601** to link to the header file **501**. Instead the data transfer engine **403** may retrieve the header file **501** by querying the blockchain **150** directly in a known fashion. E.g. a spendable output **203Y.sub.0** of the first transaction **152Y** (perhaps for a dust amount) may be addressed to Bob **103b**, thus enabling him to find it on the blockchain **150** using conventional wallet functionality. And/or Alice **103a** or a third party may signal the address of the first transaction **152Y** to Bob over an off-chain side channel (not shown). In such embodiments the hypertext engine **402** is not needed as part of the wallet function **406**.

(90) By whatever means accessed, the header file **501** is retrieved by the data transfer engine **403** of the wallet **406** and passed to the program launcher **405**. The program launcher **405** is configured to read the one or more transaction IDs **504** from the retrieved header file **501** and use these as pointers, in order to cause the data transfer engine **403** of the wallet **406** to retrieve the respective body section(s) **502** from the respective second transaction(s) **502** on the blockchain (again by querying a storage node or nodes **104S**). The program launcher **405** then combines the program header information **503** with the one or more retrieved body sections **502** in order to reconstruct the

program file **507**. The program launcher **405** may run the reconstructed program file. Alternatively or additionally it could store the reconstructed program file **507** in memory **611** or back in another transaction **152** on the chain **150**, and/or it could communicate the reconstructed file **507** to a third party for them to run or store.

(91) At least part of at least one of the body sections **502** comprises program code of the desired application. In embodiments, one or more of the body sections **502** may also comprise data to be operated on by the program code (i.e. target data or operand data), and/or one or more libraries for use in compiling and/or linking the program code.

(92) The program code of the body section(s) **502** of the program file **507** may be written in any suitable language such as C, C++, Java Script, Python, Basic, etc. It is written in a language other than the scripting language used for the locking scripts (e.g. Script). Hence this code may be referred to in embodiments as non-Script code. The program may have any application function desired by the developer, e.g. a game, office tool, communication client, or diagnostic tool, etc. The consuming party **103b** could be an end user or an operator of a node **104** such as a miner. The consuming equipment (client equipment) **102b** may comprise user equipment, server resource, or a combination thereof. In some cases the consumer equipment **102b** may additionally take the role of a node **104** of the blockchain network **106** rather than being a pure consumer.

(93) In embodiments, the program launcher **405** may take the form of a virtual machine. It may comprise a compiler **603**, assembler **604** and/or linker **605**. It may comprise one or more pre-stored libraries **606**. The running of the program file **507** may comprise compiling high-level code, assembling assembly-level code, and/or linking. In the case of linking this may comprise linking between sections **502** from different transactions, linking parts within a given section **502**, and/or linking with one of the pre-stored libraries **606**. In further alternative embodiments, the running of the program file **507** may comprise interpreting high-level code of an interpreted language. In yet another alternative, the program could already comprise low-level machine code in which case the running comprises executing the machine code.

(94) Note: it is also not excluded that the header file **501** itself could include a body section (not shown) comprising part of the program code, which may be included as one of the sections in the final reconstructed program file **507** along with the one or more body sections **502** pointed to by the transaction IDs **504**.

(95) In embodiments, the header file **501** may optionally comprise a portion **505** specifying one or more conditions for running one, more or all of the body sections **502**. These control under what circumstances the body section(s) **502** can be run. For instance these may comprise one or more access conditions determining who can access one, more or all of the body sections in order to reconstruct or run the file. More generally the conditions **505** could comprise any conditions associated with one, more or all of the body sections **502**; e.g. a compatibility condition such as a condition on what operating system the client equipment **102b** is using (e.g. to check Bob is not trying to run a windows program on a mac), or whether the client equipment **102b** has sufficient memory. Preferably the header file **501** also comprises a code stub **506** for evaluating the conditions **505**. This is not one of the body sections **502** and does not need to be included in the final reconstructed program file **507**. Instead, the code stub **506** is automatically run by the code launcher when the header file **501** is retrieved from the blockchain **150**. Like the body sections however, the code stub **506** may be written in a language other than the scripting language used for the locking scripts (e.g. a non-Script language).

(96) When run, the code stub **506** implements the relevant functionality for evaluating the conditions **505**. Alternatively it is not excluded that the evaluation is performed by pre-stored code on Bob's equipment **102b**, or even by Alice or a third party. Note also that while shown schematically as separate elements in FIG. 5, in some embodiments that do employ the code stub **506**, then the conditions **505** could be integrated into the stub **506** as an integral part of the code of the stub **506**. Alternatively however, the conditions **505** could indeed be specified in a distinct

portion of the header file **501** and referenced by the code stub **506**.

(97) The conditions **505** may for example comprise one or more access conditions, whereby access to one or more of the body sections **502** is conditional on the identity of the consuming party **103b** who is attempting to reconstruct the program file **107**. In this case the evaluation comprises checking that the identity of the consuming party **103b** is the same as an identity specified in the access conditions **505**. It may also comprise authenticating the identity of the consuming party **103b**. Suitable authentication and identity checking techniques will, in themselves, be familiar to a person skilled in the art.

(98) As another alternative or additional example, the conditions **505** may comprise one or more compatibility conditions, being conditions on the compatibility of the client computer equipment **102b** of the consuming party **103b** who is attempting to reconstruct the program file **107** (i.e. compatibility with one or more of the body sections **502**). In this case the evaluation comprises checking that the computer equipment **102b** meets one or more technical specifications specified in the compatibility conditions **505**. Suitable compatibility checking techniques will, in themselves, be familiar to a person skilled in the art.

(99) If the condition(s) **505** specified in the header file **501** is/are not met, then the code launcher **405** will either i) not access one, more or all of the body sections **502** from the blockchain **150** at all; or ii) at least will not include it/them in the final reconstructed program file **507**. This may mean no program file **507** is reconstructed, or that a version of the program file **507** with only reduced functionality is created (lacking any body sections **502** for which the consuming party **103** and/or his computer equipment **102b** did not meet the associated condition(s) **505**).

(100) In embodiments the list **504** in the header file **501** points to a plurality of transaction IDs storing a plurality of body sections for the program file **507**. In some such embodiments, there may be different respective condition **505** associated with different ones of the body portions. This can be used to create different tiers of access. An example application of this would be a program that reads from and writes to a text file. The “READ” function and “WRITE” function can be two separate parts of the code, each with their own access control.

(101) In some embodiments the header file **501** could be stored on the chain **150** in encrypted form, encrypted based on a shared secret shared between Alice and Bob. This enables only Bob (and anyone else who knows the shared secret), but not the general public, to access the header file **501** and thus reconstruct the program file **507**. The shared secret could be obtained by means of a key negotiation process conducted between Alice and Bob over an off-chain side channel (not shown), or via an on-chain process. Alternatively in other embodiments, the header file **501** need not be encrypted and may instead be made freely available to the general public via the blockchain **150**.

(102) Data stored in a transaction can remain in the UTXO set. In embodiments, to update the information, one can spend the output containing the data. The previous data is recorded on the blockchain, but only the latest information is in the UTXO set. Therefore, a blockchain pointer can always reference the data stored in the UTXO set as the latest data.

(103) A record of the most up-to-date UTXOs is stored in a UTXO set at each storage and/or mining node **104S**, **104M**. This is a reflection of the latest state of blockchain **150**, typically kept by all miners and some nodes that store an entire copy of the blockchain **150**. Once recorded on a blockchain an existing transaction cannot be modified. Instead therefore, in embodiments an update can be implemented by later adding a new transaction to the chain, wherein the new transaction spends a spendable output of the existing transaction which stores the data that is being updated. Therefore existing transaction (the one being updated) would contain at least a spendable output (and possibly an unspendable output if the data is included under an OP_RETURN or the like). Once validated, the spendable output is included in the UTXO set which is maintained at some or all of the storage and/or mining nodes **104S**, **104M**. To update, one would spend the spendable UTXO of the existing transaction and include the updated data in the new transaction (e.g. again in a new unspendable output of the new transaction using an OP_RETURN opcode or the like). Once

validated, this new transaction is now included in the UTXO set. Therefore, the latest information will always be available in the UTXO set maintained by each storage and/or mining node **104S/M**. (104) Some example implementation details of the above-described concepts are now discussed in more detail.

(105) An executable file consists of assembly-level code along with any data, libraries, or scripts that are required by the code at run-time. Typically, this is organized in a standardized format depending on the operating system. For example, portable executable (PE) format in Windows, as well as the executable and linkable format (ELF) in Linux. Although the structure varies by format, it typically consists of the following components: a program header specifying information such as machine and version information, file size, and pointer to file data; a section header table referencing the data included in the file; and a data section including the code and other requirements for execution.

(106) Specified in the file are the references to the memory addresses where the required data is stored. This file and each of its dependencies are stored on the disk and reconstructed in the memory at runtime based on the instructions specified in the file.

(107) Herein, there is disclosed an analogous executable file format where, in place of storing the file and dependencies on a local disk, it is stored in a decentralized manner on a blockchain. This file format may be referred to herein as the blockchain executable format (BEF).

(108) The assembly level code, required libraries and possibly even import data are stored, compressed and referenced on the blockchain **150**. Each component is stored in a separate transaction **152** and the header file **501** would only have to keep the TxID(s) **504** along with any header info **503**, conditions **505** and/or stub **506**. The construction, linking and execution processes can all be performed by an integrated smart wallet/virtual machine as shown in the figure below.

(109) The header file **501** has a header **503** containing some or all of the information listed above (version, checksum etc.). When fragmenting the BEF file **507**, the executable file stub **506** and header **503** are stored in a single transaction **152Y** along with the TxIDs **504** needed for reassembly and any access checks **505**. This transaction **152Y** can also be called a compressed version of the file **507**. The stub **506** is the executable code that checks the access conditions for the entire file **507**.

(110) The BEF execution would occur as follows. The method is also illustrated in FIGS. 7A to 7B.

(111) In the steps of FIG. 7A, the initial file request is initiated from a previous blockchain-embedded HTML document, user input or locally-running code, and sent to the blockchain network **106**. In response the header file **501** is returned and interpreted by the smart wallet **406**.

(112) At step **S0**, the TxID of the transaction containing the header file **501** is selected by the user **103b**. This can be done via an existing hyperlink in a hypertext document, locally-saved address, or blockchain pointer in code being executed locally.

(113) At step **S1**, the smart wallet **406** requests a BEF file via a TxID request that refers to the transaction **152Y** containing the header file **501**.

(114) At step **S2**, the network **106** responds with the transaction information and is interpreted by the smart wallet **406**.

(115) The steps of FIG. 7B show how the smart wallet checks compatibility and verification conditions **505** laid out in the header file **501**. E.g. this may be done by running the stub **506**. If all requirements are met, the wallet **406** requests the parts of the body **502** addressed in the TxID list **503** of the header file **501**, and the executable file **507** is assembled locally. The executable file, along with any local libraries required, are passed to the linker **605**, which then creates a file in memory **610** or **611** and executed on the CPU **609**. The results can be presented the user through the UI device(s) **607**, written to the blockchain **150** or used in an additional script.

(116) At step **S3**, on seeing the file format structure flag in the retrieved transaction, the wallet **406** interprets BEF file **507** and requests any further transactions referenced in the BEF document, if any (checking any access/key conditions if applicable).

(117) At step **S4**, the assembled file **507**, and any local libraries, are passed to the linker **605** in the virtual machine **405**.

(118) At step **S5**, machine code and any data are pulled from disk **611** at runtime, assembled in memory **610** and executed on the CPU **609**.

(119) At step **S6**, the output is interpreted according to the program code, which may include (but not limited to): i) display output to user **103b**, ii) push to a stack for in-Script computations, iii) feed as input into another code running in the virtual machine **405**, and/or iv) writing to the blockchain **150** via another transaction created by the smart wallet **406** as new data or to update existing data.

(120) The structure of the BEF file is shown in FIG. 5. Here is presented the expected format of a BEF file to be interpreted by the smart wallet **406**. The assembled structure **507** resembles a conventional Windows PE file but for compression, the file is deconstructed into several different transactions allowing data to be managed independently and restrict access if desired. The header file **501** contains all necessary TXID pointers **504** for reconstruction. The left-hand panel a) of FIG. 5 shows the fragmented executable file distributed across several transactions **152Y**, **152Z**. The header file **501** includes the STUB program **506**, the header details **503** and references **504** to the transactions holding the SECTIONS table and data needed for execution. The right-hand structure shows the locally assembled structure.

(121) The stub **506** can be used to control where the file **507** can be executed or under which conditions. The rest of the header file **501** may also contain code that determines the flow and format of the resulting program. This can be a simple loader or can use multiple UTXOs and transactions. The interpreted code can be segmented into different forms of software depending on user interaction or for pre-defined automata use cases. The header file **501** can have the following form. This shows the parts **503**, **504** and **505**. The stub **506** can be appended onto this format.

(122) TABLE-US-00001 Variable Type Value Unsigned Short Integer Version and System #.
Unsigned Short Integer Pages in file (# of UTXO links) Unsigned Short Integer Relocations
Unsigned Short Integer Size of Header Unsigned Short Integer Minimum Extra Sections Unsigned
Short Integer Maximum Extra Sections Unsigned Short Integer UTXO – Relocation Table
Unsigned Short Integer Overlay (#) Unsigned Short Integer Checksum Unsigned Short Integer
Reserved Sections/Extensions Long Integer Ordering information + Verification Conditions

(123) The UTXO relocation table is the list of pointers **504** to one or more sections in one or more other transactions **152Z**. In this implementation it actually identifies specific UTXOs within the referenced TxIDs, not just the TxIDs. So for example it would reference unspendable UTXO.sub.1 of TxID.sub.Z.

(124) The end of the header includes ordering information for the different sections **502**. When multiple sections **502** are to be assembled, an order should be maintained for execution as with standard compilation/execution). This may be implemented by including an assemble sequence in the header file **501** (e.g. in the header information **501**).

(125) Alternatively, that the header file **501** may only contain a pointer to a single other TxID. In this case, if there are multiple sections **501** to assemble, then each but the last section **502** will itself contain a small header portion comprising at least a pointer to the next TxID in an ordered chain. The program launcher **405** will then follow each pointer in the chain in order to collect together all the sections **502** for reconstructing the overall file **507**. In embodiments the header portion of each section **502** could comprise some or all of the same format as the header file **501**, such as the format shown in the above table. A combination of approaches could also be used. E.g. so the header file **501** contains a list **504** to some sections **502** of the overall file **507**, and at least one of those sections itself contains a header portion pointing to one or more other sections. Again the program launcher **405** follows each pointer in order to collect the sections **502** required to reconstruct the overall file **507**.

(126) For instance, the ordering information at the end of the above header format may comprise a

reference to the next part of the file that can be stored in a separate transaction and retrieved by the wallet **406** should certain verification conditions be met. This may have an error for incorrect use and can flag the need for signatures or that a UTXO following this is not of the correct type or needs permissions such as keys or fulfilment of a payment condition.

(127) This may be implemented in form of a key validation or a selective decryption function such as XOR-based decryption. Upon reconstructing the code from one or more data sources in one or various UTXOs on the blockchain, the method jumps to an OEP (original execution point) after the code has been assembled. This executable is then passed to the linker **605** and executed on the local computer. The data is loaded as an initial disk image file (such as ISO) and is reconstructed in a virtual machine or possibly a browser.

(128) The header file **501** is constructed independently in a single initiating transaction that can be read either publicly or with a selected key (such as an XOR-encrypted software link, or encoded as disclosed in PCT/IB2017/050856). This information, received and interpreted by the smart wallet **406**, allows the virtual machine **405** to determine the version of software and ensure compatibility. An example of the way it does this is explained in the next paragraph.

(129) Executing a stub program: as defined above, the smart wallet **406** incorporated into a device, a virtual machine, or browser takes the information and determines it is an executable image file via a reference flag in the transaction. The first part of the BEF file to be executed is the STUB, similar to that of the PE format. In the PE format, the DOS STUB verifies that the file is not being executed in DOS, and if it is, prints "This program cannot be run in DOS mode." and stops execution. In the context of the BEF file, the STUB can be used to specify the running environment (Windows, Linux, or a specified virtual machine) and can even be used to initiate access authentication. The client credentials can be interpreted and if insufficient information is presented for access on say a secure area of the virtual machine, the STUB program can be used to prompt the user for identity verification such as a username and password. A protocol can also be specified in the case of failure to prove access credentials such as: I) printing a message to the user and terminating execution, II) limited functionality by assembling only specified parts of the execution data, III) limited access to input data, and/or IV) blacklisting the IP address that was attempting to access the executable file.

CONCLUSION

(130) It will be appreciated that the above embodiments have been described by way of example only.

(131) More generally, according to one aspect disclosed herein there is provided a method of running a program, the method comprising, by computer equipment of a consuming party: retrieving a header file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network. The header file comprises program header information, and reference information comprising one or more respective transaction IDs of one or more respective second transactions stored on the blockchain, wherein each of the second transactions comprises a respective section of a body of the program file, at least one of the sections comprising runnable program code. The method further comprises, from the blockchain, retrieving the respective section of the program file from at least one of the one or more second transactions based on the reference information; and running a program comprising the program header information and the at least one retrieved section.

(132) In embodiments, said running of the program may comprise: reconstructing a program file comprising said program from the header information and the at least one retrieved section, storing the reconstructed program file in memory of the computer equipment of the consuming party, and running the reconstructed program file on the computer equipment of the consuming party.

(133) Alternatively the consuming party could execute each section individually and then pass on to the next section without assembling everything as a single executable file.

(134) In embodiments the code may comprise assembly-level code and the running may comprise

assembling the assembly-level code and executing the assembled code. Alternatively, in embodiments the code may comprise a high-level language and the running may comprise compiling the high-level code of the program file and executing the compiled program, e.g. using just-in-time (JIT) compilation. As another alternative, the program file may comprise raw machine code and the running may comprise executing the machine code. In yet another alternative the code may comprise a high-level language and the running may comprise interpreting high-level code. (135) In embodiments the running may further comprise linking, e.g. linking between different sections from different transactions, or between different subsections within a given section from a given transaction.

(136) In embodiments, said reference information may comprise a plurality of transaction IDs of a plurality of second transactions each storing a respective body section of the program file.

(137) The different sections may comprise program code, one or more libraries, and/or data to be operated upon by the program.

(138) In embodiments, said retrieving may comprise retrieving, from the blockchain, the respective sections from two or more of the plurality of second transactions; and the program that is run may comprise the program header information and the retrieved sections from the two or more second transactions.

(139) Alternatively the reference information in the header file could comprise only a single transaction ID referencing only a single body section in a single second transaction.

(140) In embodiments, at least one of the body sections in at least one of the second transactions may itself comprise a header portion comprising at least a reference to a third transaction containing a further body section. In this case, the program that is run may comprise the program header and the body sections from the second and third transactions.

(141) In embodiments, the header file may further comprise one or more conditions; and the method may comprise evaluating the conditions, wherein the retrieval or the running of at least one of the sections of the program is conditional on an outcome of the evaluation.

(142) In embodiments the evaluation may be performed by the computer equipment of the consuming party.

(143) In embodiments, the header file may comprise a runnable code stub for evaluating the one or more conditions. In this case the method may comprise running the code stub on the computer equipment of the consuming party in order to perform said evaluation. In embodiments, the conditions may comprise at least one access condition, being a condition on an identity of the consuming party. In this case the evaluation comprises: verifying that the identity of the consuming party matches an identity specified by the at least one access condition, wherein the retrieval or the running of at least one of the sections of the program is conditional on a positive outcome of the verification.

(144) In embodiments, the conditions may comprise at least one compatibility condition, being a condition on a compatibility of the computing equipment of the consuming party. In this case the evaluation comprises: checking that the computer equipment of the consuming party meets one or more technical criteria specified by the compatibility condition, wherein the retrieval or the running of at least one of the sections of the program is conditional on a positive outcome of the check.

(145) The technical criteria for compatibility could comprise one or more criteria placed on the hardware of the computer equipment of the consuming party. Alternatively or additionally, the technical criterion for compatibility could comprise one or more criteria placed on the software installed on the computer equipment of the consuming party, e.g. one or more criteria on the operating system, and/or one or more criteria for one or more sister applications that need to be installed.

(146) In embodiments, the conditions may comprise a respective condition associated with each of the plurality of sections, at least some of the respective conditions being different to one another. The method may comprise evaluating the conditions associated with each of at least some of the

plurality of sections, wherein the retrieval or running of each such section is conditional on the outcome of the evaluation.

(147) This advantageously provides for different tiers of access or operability, whereby some sections are made available for use and some not on a modular basis, depending on the conditions (e.g. depending on the identity of the consuming party).

(148) In embodiments the conditions may comprise a respective access condition on an identity of the consuming party for each of some or all of the sections, the evaluation comprising verifying that the identity of the consuming party matches a respective identity specified by the respective condition for each respective section, wherein the retrieval or the running of the respective section is conditional on a positive outcome of the respective verification.

(149) In embodiments the conditions may comprise a respective compatibility condition on a compatibility of the computer equipment of the consuming party with the respective section for each of some or all of the sections, the evaluation comprising verifying that the computer equipment of the consuming party meets a respective one or more technical criteria specified by the respective condition for each respective section, wherein the retrieval or the running of the respective section is conditional on a positive outcome of the respective check.

(150) In further alternative or additional embodiments, since the non-script code is more versatile than the locking/locking scripts typically used in blockchain transactions, then the one or more access conditions may comprise any user-defined condition that the developer wishes.

(151) In embodiments the program header information may comprises one or more of: a version number of the program file, a version number of system upon which the program file is to be run, a size of the program file, a description, a list of reserved variables, a date and/or time stamp, a checksum, and/or ordering information for the body sections.

(152) E.g. in embodiments the program header information comprises a minimum of the version and size information.

(153) In embodiments, the blockchain and blockchain network may be configured to operate according to an output-based model, wherein each transaction comprises at least one output and at least one input, each output comprising a locking script and each input comprising a pointer to an output of another transaction and an unlocking script for unlocking the locking script of the pointed-to output, wherein the scripts are formulated in a scripting language. In this case the header file and said program code may be formulated (written) in a language other than the scripting language.

(154) In embodiments, each of the header file and the one or more sections may be stored in at least one unspendable output of the respective first or second transaction on the blockchain.

(155) For instance, each unspendable output may be made unspendable by inclusion of an opcode of the scripting language that terminates the script when run by a node of the blockchain network. E.g. said opcode may be OP_RETURN (the scripting language being Script).

(156) In embodiments, said reference information in the header file may comprise an identifier of an individual respective output within each second transaction, and the retrieval may comprise retrieving each of the retrieved sections from the identified output.

(157) In embodiments, the header file as stored in the first transaction may be stored in encrypted form, encrypted based on a shared secret shared between the consuming party and a providing party who provides the program. In this case, the retrieval of the header file comprises decrypting the header file based on the shared secret.

(158) Alternatively the header file could be unencrypted and thus freely available on the chain to the public.

(159) According to another aspect disclosed herein, there is provided a computer program embodied on computer-readable storage and configured so as when run on the computer equipment of the consuming party to perform any of the method steps disclosed herein.

(160) According to another aspect, there is provided the equipment of the consuming party,

comprising: memory comprising one or more memory units, and processing apparatus comprising one or more processing units; wherein the memory stores code arranged to run on the processing apparatus, the code being configured so as when run to perform the method of the consuming party according to any embodiment disclosed herein.

(161) According to another aspect disclosed herein, there is provided a group of transactions for recordal in a blockchain, the group comprising, embodied on a computer-readable data medium or media: a first transaction comprising a header file, and one or more second transactions each comprising a respective section of a body of the program file, at least one of the sections comprising runnable program code; wherein the header file comprises program header information and a set of one or more respective transaction IDs of one or more of the second transactions stored on the blockchain, enabling a program file to be constructed from the program header information and at least one of the sections.

(162) According to another aspect disclosed herein, there may be provided a method comprising the actions of the first party, second party, any third party that may be involved, and/or the network of nodes.

(163) According to another aspect disclosed herein, there may be provided a system comprising the computer equipment of the first party, the computer equipment of the second party, the computer equipment of any third party, and/or the network of nodes.

(164) Other variants or use cases of the disclosed techniques may become apparent to the person skilled in the art once given the disclosure herein. The scope of the disclosure is not limited by the described embodiments but only by the accompanying claims.

Claims

1. A method of running a program, the method comprising, by computer equipment of a consuming party: retrieving a header file of a program file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network, the header file comprising: program header information, and reference information comprising one or more respective transaction identifiers (IDs) of one or more respective second transactions stored on the blockchain, wherein each respective one of the second transactions comprises a respective section of respective sections of a body of the program file, at least one respective section of the respective sections comprising runnable program code; from the blockchain, retrieving the respective section of the program file from at least one of the one or more respective second transactions based on the reference information; and running a program comprising the program header information and the at least one retrieved section.

2. The method of claim 1, wherein said running of the program comprises: reconstructing a program file comprising said program from the header information and the at least one retrieved section, storing the reconstructed program file in memory of the computer equipment of the consuming party, and running the reconstructed program file on the computer equipment of the consuming party.

3. The method of claim 1, wherein said reference information comprises a plurality of transaction IDs of a plurality of second transactions each storing a respective body section of the program file.

4. The method of claim 3, wherein: said retrieving of the respective sections comprises retrieving, from the blockchain, the respective sections from two or more of the plurality of second transactions; and the program that is run comprises the program header information and the retrieved sections from the two or more second transactions.

5. The method of claim 1, wherein: at least one of the body sections in at least one of the second transactions itself comprises a header portion comprising at least a reference to a third transaction containing a further body section; and the program that is run comprises the program header and the body sections from the second and third transactions.

6. The method of claim 1, wherein the header file further comprises one or more conditions; and the method comprises evaluating the conditions, wherein the retrieval or the running of at least one of the sections of the program is conditional on an outcome of the evaluation.
7. The method of claim 6, wherein the evaluation is performed by the computer equipment of the consuming party.
8. The method of claim 7, wherein the header file comprises a runnable code stub for evaluating the one or more conditions, the method comprising running the code stub on the computer equipment of the consuming party in order to perform the evaluation.
9. The method of claim 6, wherein the conditions comprise at least one access condition, being a condition on an identity of the consuming party; the evaluation comprising verifying that the identity of the consuming party matches an identity specified by the at least one access condition, wherein the retrieval or the running of at least one of the sections of the program is conditional on a positive outcome of the verification.
10. The method of claim 6, wherein the conditions comprise at least one compatibility condition, being a condition on a compatibility of the computing equipment of the consuming party; the evaluation comprising checking that the computer equipment of the consuming party meets one or more technical criteria specified by the compatibility condition, wherein the retrieval or the running of at least one of the sections of the program is conditional on a positive outcome of the checking method step.
11. The method of claim 6, wherein: said reference information comprises a plurality of transaction IDs of a plurality of second transactions each storing a respective body section of the program file; the conditions comprise a respective condition associated with each of the plurality of respective sections, at least some of the respective conditions being different to one another; and the method comprises evaluating the conditions associated with each of at least some of the plurality of sections, wherein the retrieval or running of each such section is conditional on the outcome of the evaluation.
12. The method of claim 11, wherein the conditions comprise a respective access condition on an identity of the consuming party for each of some or all of the sections, the evaluation comprising verifying that the identity of the consuming party matches a respective identity specified by the respective condition for each respective section, wherein the retrieval or the running of the respective section is conditional on a positive outcome of the respective verification.
13. The method of claim 6, wherein: wherein said reference information comprises a plurality of transaction IDs of a plurality of second transactions each storing a respective body section of the program file; and the conditions comprise a respective compatibility condition on a compatibility of the computer equipment of the consuming party with the respective section for each of some or all of the sections, the evaluation comprising verifying that the computer equipment of the consuming party meets a respective one or more technical criteria specified by the respective condition for each respective section, wherein the retrieval or the running of the respective section is conditional on a positive outcome of the respective check.
14. The method of claim 1, wherein the program header information comprises one or more of: a version number of the program file, a version number of system upon which the program file is to be run, a size of the program file, a description, a list of reserved variables, a date and/or time stamp, a checksum, and/or ordering information for the body sections.
15. The method of claim 1, wherein the blockchain and blockchain network are configured to operate according to an output-based model, wherein each transaction comprises at least one output and at least one input, each output comprising a locking script and each input comprising a pointer to an output of another transaction and an unlocking script for unlocking the locking script of the pointed-to output; wherein the locking and unlocking scripts are formulated in a scripting language, and wherein the header file and said program code is formulated in a language other than the scripting language.

16. The method of claim 15, wherein each of the header file and the one or more sections is stored in at least one unspendable output of the respective first or second transaction on the blockchain.

17. The method of claim 15 or 16, wherein said reference information in the header file comprises an identifier of an individual respective output within each second transaction, and the retrieval comprises retrieving each of the retrieved sections from the identified output.

18. The method of claim 1, wherein the header file as stored in the first transaction is stored in encrypted form, encrypted based on a shared secret shared between the consuming party and a providing party who provides the program; and wherein the retrieval of the header file comprises decrypting the header file based on the shared secret.

19. A computer program embodied on computer-readable storage and configured so as when run on computer equipment of a consuming party to perform a method of: retrieving a header file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network, the header file comprising: program header information, and reference information comprising one or more respective transaction identifiers of one or more respective second transactions stored on the blockchain, wherein each respective one of the second transactions comprises a respective section of respective sections of a body of a program file, at least one respective section of the respective sections comprising runnable program code; from the blockchain, retrieving the respective section of the program file from at least one of the one or more respective second transactions based on the reference information; and running a first program comprising the program header information and the at least one retrieved section.

20. Equipment of a consuming party, comprising: memory comprising one or more memory units, and processing apparatus comprising one or more processing units; wherein the memory stores code arranged to run on the processing apparatus, the code being configured so as when run to perform a method of: retrieving a header file from a first transaction recorded on a blockchain maintained across a plurality of nodes of a blockchain network, the header file comprising: program header information, and reference information comprising one or more respective transaction identifiers of one or more respective second transactions stored on the blockchain, wherein each respective one of the second transactions comprises a respective section of respective sections of a body of a program file, at least one respective section of the respective sections comprising runnable program code; from the blockchain, retrieving the respective section of the program file from at least one of the one or more respective second transactions based on the reference information; and running a program comprising the program header information and the at least one retrieved section.
