

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411962
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Marcon; Paulo et al.

Managed run-time environment-based detection of a ransomware attack

Abstract

An illustrative method includes a data protection system detecting an anomaly associated with a managed run-time environment; determining, based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; and performing, based on the determining that the managed run-time environment is possibly being targeted by the security threat, a remedial action associated with the managed run-time environment. The performing the remedial action may include identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity.

Inventors: Marcon; Paulo (Porto Alegre, BR), Gelinger Maciel; Romulo (Canoas, BR), Karr; Ronald (Palo Alto, CA)

Applicant: Pure Storage, Inc. (Santa Clara, CA)

Family ID: 89576496

Assignee: Pure Storage, Inc. (Santa Clara, CA)

Appl. No.: 18/233686

Filed: August 14, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240028748 A1	Jan. 25, 2024

Related U.S. Application Data

continuation parent-doc US 17161553 20210128 US 11520907 20221206 child-doc US 17980354
continuation parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US 17725182
continuation-in-part parent-doc US 18127926 20230329 US 12204657 child-doc US 18233686
continuation-in-part parent-doc US 17980354 20221103 US 11720691 20230808 child-doc US
18127926
continuation-in-part parent-doc US 17846301 20220622 US 12248566 child-doc US 18127926
continuation-in-part parent-doc US 17725182 20220420 US 11657146 20230523 child-doc US
18127926
continuation-in-part parent-doc US 17723903 20220419 US 12079502 child-doc US 18127926
continuation-in-part parent-doc US 17541870 20211203 US 12153670 child-doc US 18127926
continuation-in-part parent-doc US 17506509 20211020 US 12079333 child-doc US 18127926
continuation-in-part parent-doc US 17506501 20211020 US 12050683 child-doc US 18127926
continuation-in-part parent-doc US 17463088 20210831 US 12067118 child-doc US 18127926
continuation-in-part parent-doc US 17409124 20210823 US 12079356 child-doc US 18127926
continuation-in-part parent-doc US 17409130 20210823 PENDING child-doc US 18127926
continuation-in-part parent-doc US 17409135 20210823 US 12050689 child-doc US 18127926
continuation-in-part parent-doc US 17342203 20210608 US 11657155 20230523 child-doc US
18127926
continuation-in-part parent-doc US 17235737 20210420 US 11687418 20230627 child-doc US
18127926
continuation-in-part parent-doc US 17074313 20201019 PENDING child-doc US 18127926
continuation-in-part parent-doc US 17039556 20200930 US 11720714 20230808 child-doc US
18127926
continuation-in-part parent-doc US 17039536 20200930 US 11625481 20230411 child-doc US
18127926
continuation-in-part parent-doc US 17039604 20200930 US 11651075 20230516 child-doc US
18127926
continuation-in-part parent-doc US 17039486 20200930 US 11720692 20230808 child-doc US
18127926
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17723903
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17409135
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17463088
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17541870
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17409124
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17342203
continuation-in-part parent-doc US 16917030 20200630 US 11675898 20230613 child-doc US
17161553
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17409130
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17235737
continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US
17506509

continuation-in-part parent-doc US 16916903 20200630 US 11341236 20220524 child-doc US 17506501
continuation-in-part parent-doc US 16711060 20191211 PENDING child-doc US 17039556
continuation-in-part parent-doc US 16711060 20191211 PENDING child-doc US 16917030
us-provisional-application US 62985229 20200304
us-provisional-application US 62939518 20191122

Publication Classification

Int. Cl.: G06F21/60 (20130101); G06F3/06 (20060101)

U.S. Cl.:

CPC G06F21/602 (20130101); G06F3/0608 (20130101); G06F3/0652 (20130101); G06F3/067 (20130101);

Field of Classification Search

CPC: G06F (21/602); G06F (3/0608); G06F (3/0652); G06F (3/067)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5208813	12/1992	Stallmo	N/A	N/A
5403639	12/1994	Belsan et al.	N/A	N/A
5940838	12/1998	Schmuck et al.	N/A	N/A
6263350	12/2000	Wollrath et al.	N/A	N/A
6412045	12/2001	DeKoning et al.	N/A	N/A
6595035	12/2002	Maley	N/A	N/A
6718448	12/2003	Ofer	N/A	N/A
6757769	12/2003	Ofer	N/A	N/A
6799283	12/2003	Tamai et al.	N/A	N/A
6834298	12/2003	Singer et al.	N/A	N/A
6850938	12/2004	Sadjadi	N/A	N/A
6915434	12/2004	Kuroda et al.	N/A	N/A
6973549	12/2004	Testardi	N/A	N/A
7028216	12/2005	Aizawa et al.	N/A	N/A
7028218	12/2005	Schwarm et al.	N/A	N/A
7039827	12/2005	Meyer et al.	N/A	N/A
7043759	12/2005	Kaashoek et al.	N/A	N/A
7149858	12/2005	Kiselev	N/A	N/A
7213264	12/2006	Poletto et al.	N/A	N/A
7216164	12/2006	Whitmore et al.	N/A	N/A
7231050	12/2006	Harris	N/A	N/A
7418730	12/2007	Chu et al.	N/A	N/A
7467333	12/2007	Keeton et al.	N/A	N/A
7559088	12/2008	Cohen et al.	N/A	N/A
7735136	12/2009	Mantripragada et al.	N/A	N/A

7739211	12/2009	Coffman et al.	N/A	N/A
7774848	12/2009	D'Mello et al.	N/A	N/A
7783666	12/2009	Zhuge et al.	N/A	N/A
7783682	12/2009	Patterson	N/A	N/A
7797436	12/2009	Banerjee et al.	N/A	N/A
7797748	12/2009	Zheng et al.	N/A	N/A
7873619	12/2010	Faibish et al.	N/A	N/A
7913300	12/2010	Flank et al.	N/A	N/A
7930260	12/2010	Deo et al.	N/A	N/A
7933936	12/2010	Aggarwal et al.	N/A	N/A
7958159	12/2010	Tran et al.	N/A	N/A
7975115	12/2010	Wayda et al.	N/A	N/A
7979613	12/2010	Zohar et al.	N/A	N/A
7996547	12/2010	Sudhakar	N/A	N/A
8001601	12/2010	Duffield et al.	N/A	N/A
8010519	12/2010	Hsu et al.	N/A	N/A
8010667	12/2010	Zhang et al.	N/A	N/A
8086652	12/2010	Bisson et al.	N/A	N/A
8117464	12/2011	Kogelnik	N/A	N/A
8132260	12/2011	Mayer et al.	N/A	N/A
8137199	12/2011	Douceur et al.	N/A	N/A
8200887	12/2011	Bennett	N/A	N/A
8205065	12/2011	Matze	N/A	N/A
8260792	12/2011	Hsu	N/A	N/A
8266267	12/2011	Guruswamy et al.	N/A	N/A
8352540	12/2012	Anglin	709/201	G06F 3/067
8356001	12/2012	Collazo	706/14	G06F 21/577
8433869	12/2012	Natanzon	707/634	G06F 11/2069
8468604	12/2012	Claudatos	726/24	G06F 21/568
8484163	12/2012	Yucel	370/216	G06F 16/21
8495037	12/2012	Westenberg	707/698	G06F 21/56
8495472	12/2012	Magerramov	714/774	G06Q 40/02
8495740	12/2012	Rouland	726/25	H04L 63/1433
8499349	12/2012	Cruz	726/30	G06F 21/568
8504797	12/2012	Mimatsu	709/213	G06F 3/0665
8504876	12/2012	Mork	714/25	G06F 21/55
8527544	12/2012	Colgrove	707/791	G06F 3/0641
8554703	12/2012	Lin	706/12	G06N 20/00
8560747	12/2012	Tan	710/200	G06F 9/526
8621241	12/2012	Stephenson	713/193	H04L 67/56
8635498	12/2013	Kahana	714/38.1	G06F 11/3476

8667583	12/2013	Polyakov	726/22	G06F 21/552
8683591	12/2013	Wittenschlaeger	726/23	H04L 63/18
8688620	12/2013	Viswanathan	706/62	G06F 11/076
8700875	12/2013	Barron	711/170	G06F 3/067
8706914	12/2013	Duchesneau	709/250	F28D 15/06
8719936	12/2013	Moffie	726/23	G06F 21/552
8745004	12/2013	Natanzon	707/634	G06F 16/10
8751463	12/2013	Chamness	711/170	G06F 16/1744
8793790	12/2013	Khurana	726/1	G06N 5/045
8806160	12/2013	Colgrove	707/639	G06F 3/061
8822155	12/2013	Sukumar	435/6.12	C12Q 1/6886
8862727	12/2013	Jayachandran et al.	N/A	N/A
8862728	12/2013	Jayachandran	709/224	H04L 41/0613
8863293	12/2013	Christodorescu	726/25	G06F 21/552
8874850	12/2013	Goodson	711/133	G06F 12/0893
8904092	12/2013	Tucek	711/206	G06F 12/0238
8918478	12/2013	Ozzie	714/763	G06F 11/1004
8959305	12/2014	Lecrone	707/813	G06F 12/023
9015426	12/2014	Stabrawa	711/154	G06F 3/0659
9032160	12/2014	Natanzon	711/135	G06F 1/24
9069955	12/2014	Dolph	N/A	H04L 63/14
9081713	12/2014	Bennett	N/A	G11C 16/10
9088601	12/2014	Friedrichs	N/A	G06F 21/564
9092146	12/2014	Jayaraman	N/A	G06F 3/0655
9111232	12/2014	Chi	N/A	G06N 5/02
9112895	12/2014	Lin	N/A	H04L 63/1425
9189334	12/2014	Bennett	N/A	N/A
9197514	12/2014	Finkelstein	N/A	H04L 41/5025
9197658	12/2014	Wittenschlaeger	N/A	N/A
9235481	12/2015	Natanzon et al.	N/A	N/A
9246935	12/2015	Lietz	N/A	H04L 67/34
9256736	12/2015	Tao	N/A	H04L 63/1466

9256746	12/2015	Toback	N/A	G06F 21/577
9258321	12/2015	Amsler	N/A	H04L 63/1425
9268797	12/2015	Whitehead	N/A	G06F 11/1456
9275063	12/2015	Natanzon	N/A	G06F 3/0611
9280678	12/2015	Redberg	N/A	G06F 16/164
9311182	12/2015	Bennett	N/A	N/A
9317686	12/2015	Ye	N/A	G06F 21/568
9324022	12/2015	Williams, Jr.	N/A	G06N 3/045
9338181	12/2015	Burns	N/A	H04L 63/1433
9348742	12/2015	Brezinski	N/A	G06F 21/566
9349103	12/2015	Eberhardt, III	N/A	G06N 20/00
9350750	12/2015	Aval	N/A	H04L 63/1425
9363286	12/2015	Golovanov	N/A	H04L 63/1475
9386009	12/2015	Marion	N/A	H04L 63/083
9395922	12/2015	Nishikido	N/A	G06F 3/065
9407646	12/2015	Cruz Mota	N/A	H04L 63/1458
9423967	12/2015	Colgrove	N/A	G06F 3/0653
9430156	12/2015	Shilane	N/A	G06F 3/067
9436396	12/2015	Colgrove	N/A	G06F 3/0685
9436720	12/2015	Colgrove	N/A	G06F 3/0673
9454476	12/2015	Colgrove et al.	N/A	N/A
9454477	12/2015	Colgrove	N/A	G06F 3/0638
9477693	12/2015	Bachu	N/A	G06F 11/1448
9479518	12/2015	Fang	N/A	H04L 63/1425
9483189	12/2015	Brooker	N/A	G06F 12/0811
9503469	12/2015	Lin	N/A	N/A
9513820	12/2015	Shalev	N/A	G06F 3/0611
9516016	12/2015	Colgrove	N/A	H04L 63/0838

9516053	12/2015	Muddu	N/A	G06F 3/04847
9552248	12/2016	Miller	N/A	G06F 11/1451
9552299	12/2016	Stalzer	N/A	G06F 12/0866
9588847	12/2016	Natanzon	N/A	G06F 11/1458
9591010	12/2016	Muddu et al.	N/A	N/A
9596254	12/2016	Muddu et al.	N/A	N/A
9609009	12/2016	Muddu et al.	N/A	N/A
9609011	12/2016	Muddu et al.	N/A	N/A
9628507	12/2016	Haq	N/A	G06F 21/564
9632870	12/2016	Bennett	N/A	G06F 3/0689
9641544	12/2016	Treat	N/A	H04L 43/04
9646262	12/2016	Phillipps	N/A	G06N 20/10
9665302	12/2016	Huff	N/A	G06F 3/0608
9665303	12/2016	Huff	N/A	G06F 3/0664
9667641	12/2016	Muddu et al.	N/A	N/A
9679165	12/2016	Pittelko	N/A	H04L 63/0428
9684472	12/2016	Bish	N/A	G06F 3/0683
9686294	12/2016	Kantor	N/A	H04W 4/40
9692785	12/2016	Stephan	N/A	H04L 63/1466
9692789	12/2016	Kirti	N/A	H04L 63/1425
9696939	12/2016	Frank	N/A	G06F 11/1471
9699205	12/2016	Muddu et al.	N/A	N/A
9705914	12/2016	Di Pietro	N/A	H04W 12/125
9705970	12/2016	Pomerantz	N/A	H04L 47/70
9716723	12/2016	Wittenschlaeger	N/A	N/A
9729572	12/2016	Adams	N/A	G06F 21/52
9736243	12/2016	Chablani	N/A	G06F 3/067
9740403	12/2016	Storer	N/A	G06F 3/067
9798754	12/2016	Shilane	N/A	G06F 11/3037
9798891	12/2016	Deulgaonkar	N/A	G06F 3/0659
9813432	12/2016	Dasgupta	N/A	G06F 11/3409
9813435	12/2016	Muddu et al.	N/A	N/A

9813500	12/2016	Mason	N/A	H04L 67/1097
9824609	12/2016	Sadeh-Konieczpol	N/A	H04L 63/1425
9838409	12/2016	Flacher	N/A	H04L 63/1425
9838410	12/2016	Muddu et al.	N/A	N/A
9838419	12/2016	Fleyder	N/A	H04L 63/101
9864874	12/2017	Shanbhag	N/A	H04L 63/10
9866573	12/2017	Kozlovsky	N/A	H04L 63/1416
9870251	12/2017	Baracaldo Angel	N/A	G06F 9/4843
9875355	12/2017	Williams	N/A	H04L 63/1408
9888032	12/2017	Dekel	N/A	G06F 21/60
9900332	12/2017	Muddu et al.	N/A	N/A
9904587	12/2017	Potlapally	N/A	G06F 11/079
9916849	12/2017	Mader	N/A	G11B 7/126
9917854	12/2017	Natanzon	N/A	G06F 11/1458
9923910	12/2017	Di Pietro	N/A	H04L 63/1416
9928155	12/2017	Xu	N/A	G06F 11/3612
9930061	12/2017	Zandani	N/A	H04L 63/1425
9940206	12/2017	Bogdanov	N/A	G06F 11/2002
9948663	12/2017	Wang	N/A	H04L 63/1416
9954884	12/2017	Hassell	N/A	H04L 63/1441
10003605	12/2017	Muddu et al.	N/A	N/A
10015177	12/2017	Muddu et al.	N/A	N/A
10025583	12/2017	Butler et al.	N/A	N/A
10025673	12/2017	Maccanti et al.	N/A	N/A
10032033	12/2017	Gu et al.	N/A	N/A
10037779	12/2017	Mendonsa et al.	N/A	N/A
10038707	12/2017	Muddu et al.	N/A	N/A
10042697	12/2017	Ahad	N/A	N/A
10078473	12/2017	Ahmed et al.	N/A	N/A
10102356	12/2017	Sahin et al.	N/A	N/A
10120887	12/2017	Patel et al.	N/A	N/A
10122740	12/2017	Finkelshtein et al.	N/A	N/A
10122752	12/2017	Soman et al.	N/A	N/A
10169601	12/2018	Arasan et al.	N/A	N/A
10185495	12/2018	Katsuki	N/A	N/A

10324639	12/2018	Seo	N/A	N/A
10359939	12/2018	Wei et al.	N/A	N/A
10409986	12/2018	Natanzon et al.	N/A	N/A
10423541	12/2018	Levin et al.	N/A	N/A
10438000	12/2018	Gu et al.	N/A	N/A
10466925	12/2018	Blanco et al.	N/A	N/A
10476907	12/2018	Hittel et al.	N/A	N/A
10496487	12/2018	Natanzon et al.	N/A	N/A
10503427	12/2018	Botes	N/A	G06F 11/1662
10503611	12/2018	Srivastav et al.	N/A	N/A
10503897	12/2018	Striem-Amit	N/A	N/A
10503904	12/2018	Singh et al.	N/A	N/A
10516688	12/2018	Tamir et al.	N/A	N/A
10530791	12/2019	Astigarraga et al.	N/A	N/A
10536482	12/2019	Gabaev et al.	N/A	N/A
10542021	12/2019	Mehr	N/A	N/A
10558809	12/2019	Joyce et al.	N/A	N/A
10567406	12/2019	Astigarraga et al.	N/A	N/A
10567500	12/2019	Leshinsky et al.	N/A	N/A
10587484	12/2019	Nazar et al.	N/A	N/A
10587632	12/2019	Perumalla et al.	N/A	N/A
10609066	12/2019	Nossik et al.	N/A	N/A
10614430	12/2019	Armstrong	N/A	N/A
10628585	12/2019	Tamir et al.	N/A	N/A
10628586	12/2019	Jung	N/A	N/A
10673628	12/2019	Shtop et al.	N/A	N/A
10685293	12/2019	Heimann et al.	N/A	N/A
10725965	12/2019	Rokicki	N/A	N/A
10733290	12/2019	Berler et al.	N/A	N/A
10735448	12/2019	Kesin et al.	N/A	N/A
10742682	12/2019	Yu et al.	N/A	N/A
10778714	12/2019	Shi	N/A	N/A
10789363	12/2019	Gursoy et al.	N/A	N/A
10810088	12/2019	Gu et al.	N/A	N/A
10819721	12/2019	Jenkins	N/A	N/A
10826918	12/2019	Rasovic	N/A	N/A
10846137	12/2019	Vallala et al.	N/A	N/A
10853200	12/2019	Chen et al.	N/A	N/A
10877683	12/2019	Wu et al.	N/A	N/A
10884914	12/2020	Ahmed et al.	N/A	N/A
10902929	12/2020	Pawlowski	N/A	N/A
10965708	12/2020	Kolychev et al.	N/A	N/A
10983877	12/2020	Hoffman	N/A	G06F 11/3034
10992581	12/2020	Kuang et al.	N/A	N/A
11005936	12/2020	Liu et al.	N/A	N/A
11010233	12/2020	Golden et al.	N/A	N/A
11025665	12/2020	Newman et al.	N/A	N/A
11030314	12/2020	Kucherov et al.	N/A	N/A

11036392	12/2020	Resch	N/A	N/A
11038902	12/2020	Zaw	N/A	N/A
11044263	12/2020	McLean et al.	N/A	N/A
11049039	12/2020	Zimmer et al.	N/A	N/A
11055411	12/2020	Strogov et al.	N/A	N/A
11055652	12/2020	Kannan et al.	N/A	N/A
11100232	12/2020	Juncker et al.	N/A	N/A
11106810	12/2020	Natanzon et al.	N/A	N/A
11170107	12/2020	Selvaraj et al.	N/A	N/A
11212681	12/2020	Balaramn et al.	N/A	N/A
11336685	12/2021	Rodriguez et al.	N/A	N/A
11494216	12/2021	Halcrow	N/A	G06F 21/554
11663340	12/2022	Wu	726/25	G06F 11/1451
2002/0035683	12/2001	Kaashoek et al.	N/A	N/A
2002/0038436	12/2001	Suzuki	N/A	N/A
2002/0087544	12/2001	Selkirk et al.	N/A	N/A
2002/0178335	12/2001	Selkirk et al.	N/A	N/A
2002/0184362	12/2001	Banerjee et al.	N/A	N/A
2003/0140209	12/2002	Testardi	N/A	N/A
2004/0049572	12/2003	Yamamoto et al.	N/A	N/A
2004/0117640	12/2003	Chu et al.	N/A	N/A
2004/0139340	12/2003	Johnson et al.	N/A	N/A
2005/0020945	12/2004	Tosaya et al.	N/A	N/A
2005/0066095	12/2004	Mullick et al.	N/A	N/A
2005/0193236	12/2004	Stager et al.	N/A	N/A
2005/0203921	12/2004	Newman et al.	N/A	N/A
2005/0216535	12/2004	Saika et al.	N/A	N/A
2005/0223154	12/2004	Uemura	N/A	N/A
2005/0223239	12/2004	Dotan	N/A	N/A
2006/0005227	12/2005	Samuelsson et al.	N/A	N/A
2006/0021029	12/2005	Brickell et al.	N/A	N/A
2006/0074940	12/2005	Craft et al.	N/A	N/A
2006/0136365	12/2005	Kedem et al.	N/A	N/A
2006/0155946	12/2005	Ji	N/A	N/A
2007/0006319	12/2006	Fitzgibbon et al.	N/A	N/A
2007/0067585	12/2006	Ueda et al.	N/A	N/A
2007/0105589	12/2006	Lu	N/A	N/A
2007/0109856	12/2006	Pellicone et al.	N/A	N/A
2007/0162954	12/2006	Pela	N/A	N/A
2007/0171562	12/2006	Maejima et al.	N/A	N/A
2007/0174673	12/2006	Kawaguchi et al.	N/A	N/A
2007/0220068	12/2006	Thompson et al.	N/A	N/A
2007/0220313	12/2006	Katsuragi et al.	N/A	N/A
2007/0245090	12/2006	King et al.	N/A	N/A
2007/0266179	12/2006	Chavan et al.	N/A	N/A
2008/0022378	12/2007	Repasi et al.	N/A	N/A
2008/0047013	12/2007	Claudatos et al.	N/A	N/A
2008/0059699	12/2007	Kubo et al.	N/A	N/A

2008/0065852	12/2007	Moore et al.	N/A	N/A
2008/0095182	12/2007	Elzur et al.	N/A	N/A
2008/0098476	12/2007	Syversen	N/A	N/A
2008/0134174	12/2007	Sheu et al.	N/A	N/A
2008/0155191	12/2007	Anderson et al.	N/A	N/A
2008/0178040	12/2007	Kobayashi	N/A	N/A
2008/0209096	12/2007	Lin et al.	N/A	N/A
2008/0244205	12/2007	Amano et al.	N/A	N/A
2008/0275928	12/2007	Shuster	N/A	N/A
2008/0282045	12/2007	Biswas et al.	N/A	N/A
2008/0285083	12/2007	Aonuma	N/A	N/A
2008/0307270	12/2007	Li	N/A	N/A
2009/0003131	12/2008	Meyer	N/A	N/A
2009/0006587	12/2008	Richter	N/A	N/A
2009/0037662	12/2008	Frese et al.	N/A	N/A
2009/0106318	12/2008	Mantripragada et al.	N/A	N/A
2009/0133122	12/2008	Koo et al.	N/A	N/A
2009/0158432	12/2008	Zheng et al.	N/A	N/A
2009/0204700	12/2008	Sudhakar	N/A	N/A
2009/0204858	12/2008	Kawaba	N/A	N/A
2009/0228648	12/2008	Wack	N/A	N/A
2009/0300084	12/2008	Whitehouse	N/A	N/A
2010/0005173	12/2009	Baskaran et al.	N/A	N/A
2010/0017870	12/2009	Kargupta	N/A	N/A
2010/0031361	12/2009	Shukla	N/A	N/A
2010/0057673	12/2009	Savov	N/A	N/A
2010/0058026	12/2009	Heil et al.	N/A	N/A
2010/0058122	12/2009	Compton et al.	N/A	N/A
2010/0067706	12/2009	Anan et al.	N/A	N/A
2010/0077205	12/2009	Ekstrom et al.	N/A	N/A
2010/0082879	12/2009	McKean et al.	N/A	N/A
2010/0106905	12/2009	Kurashige et al.	N/A	N/A
2010/0153620	12/2009	McKean et al.	N/A	N/A
2010/0153641	12/2009	Jagadish et al.	N/A	N/A
2010/0191897	12/2009	Zhang et al.	N/A	N/A
2010/0250713	12/2009	Sudhakar	N/A	N/A
2010/0250802	12/2009	Waugh et al.	N/A	N/A
2010/0250882	12/2009	Hutchison et al.	N/A	N/A
2010/0281225	12/2009	Chen et al.	N/A	N/A
2010/0287327	12/2009	Li et al.	N/A	N/A
2010/0306849	12/2009	Zheng et al.	N/A	N/A
2011/0035540	12/2010	Fitzgerald et al.	N/A	N/A
2011/0072300	12/2010	Rousseau	N/A	N/A
2011/0072515	12/2010	Park et al.	N/A	N/A
2011/0121231	12/2010	Tian et al.	N/A	N/A
2011/0125716	12/2010	Drews et al.	N/A	N/A
2011/0145497	12/2010	Bish et al.	N/A	N/A
2011/0145598	12/2010	Smith et al.	N/A	N/A
2011/0161559	12/2010	Yurzola et al.	N/A	N/A
2011/0167221	12/2010	Pangal et al.	N/A	N/A

2011/0191537	12/2010	Kawaguchi et al.	N/A	N/A
2011/0238634	12/2010	Kobara	N/A	N/A
2011/0258164	12/2010	Mandagere et al.	N/A	N/A
2011/0271146	12/2010	Mork et al.	N/A	N/A
2011/0296237	12/2010	Mandagere et al.	N/A	N/A
2012/0023375	12/2011	Dutta et al.	N/A	N/A
2012/0036309	12/2011	Dillow et al.	N/A	N/A
2012/0096225	12/2011	Khawand et al.	N/A	N/A
2012/0117029	12/2011	Gold	N/A	N/A
2012/0151593	12/2011	Kang et al.	N/A	N/A
2012/0159629	12/2011	Lee et al.	N/A	N/A
2012/0191933	12/2011	Zbiciak et al.	N/A	N/A
2012/0198175	12/2011	Atkisson	N/A	N/A
2012/0198192	12/2011	Balasubramanian et al.	N/A	N/A
2012/0221803	12/2011	Stabrawa et al.	N/A	N/A
2012/0296922	12/2011	Hsu	N/A	N/A
2012/0303868	12/2011	Tucek	N/A	N/A
2012/0311708	12/2011	Agarwal et al.	N/A	N/A
2012/0324236	12/2011	Srivastava et al.	N/A	N/A
2012/0330954	12/2011	Sivasubramanian et al.	N/A	N/A
2013/0036465	12/2012	Chuan et al.	N/A	N/A
2013/0042052	12/2012	Colgrove et al.	N/A	N/A
2013/0046995	12/2012	Movshovitz	N/A	N/A
2013/0047029	12/2012	Ikeuchi et al.	N/A	N/A
2013/0061322	12/2012	Sethumadhavan et al.	N/A	N/A
2013/0091102	12/2012	Nayak	N/A	N/A
2013/0117507	12/2012	Chiang	N/A	N/A
2013/0205110	12/2012	Kettner	N/A	N/A
2013/0227236	12/2012	Flynn et al.	N/A	N/A
2013/0254412	12/2012	Menezes et al.	N/A	N/A
2013/0275391	12/2012	Batwara et al.	N/A	N/A
2013/0275656	12/2012	Talagala et al.	N/A	N/A
2013/0283058	12/2012	Fiske et al.	N/A	N/A
2013/0290648	12/2012	Shao et al.	N/A	N/A
2013/0318314	12/2012	Markus et al.	N/A	N/A
2013/0339303	12/2012	Potter et al.	N/A	N/A
2013/0340082	12/2012	Shanley	N/A	N/A
2014/0004829	12/2013	Rieger et al.	N/A	N/A
2014/0006760	12/2013	Nemiroff et al.	N/A	N/A
2014/0052946	12/2013	Kimmel	N/A	N/A
2014/0059688	12/2013	Margalit	N/A	N/A
2014/0068791	12/2013	Resch	N/A	N/A
2014/0089730	12/2013	Watanabe et al.	N/A	N/A
2014/0101361	12/2013	Gschwind	N/A	N/A
2014/0143517	12/2013	Jin et al.	N/A	N/A
2014/0172929	12/2013	Sedayao et al.	N/A	N/A
2014/0196115	12/2013	Pelykh	N/A	N/A
2014/0201150	12/2013	Kumarasamy et al.	N/A	N/A

2014/0201836	12/2013	Amsler	N/A	N/A
2014/0215129	12/2013	Kuzmin et al.	N/A	N/A
2014/0229131	12/2013	Cohen et al.	N/A	N/A
2014/0229452	12/2013	Serita et al.	N/A	N/A
2014/0230051	12/2013	Vallinayagam et al.	N/A	N/A
2014/0230062	12/2013	Kumaran	N/A	N/A
2014/0236891	12/2013	Talius et al.	N/A	N/A
2014/0281155	12/2013	Cohen	N/A	N/A
2014/0281308	12/2013	Lango et al.	N/A	N/A
2014/0283061	12/2013	Quinlan et al.	N/A	N/A
2014/0289853	12/2013	Teddy et al.	N/A	N/A
2014/0325115	12/2013	Ramsundar et al.	N/A	N/A
2015/0047045	12/2014	Salsamendi	N/A	N/A
2015/0081950	12/2014	Nezu	N/A	N/A
2015/0106557	12/2014	Yu et al.	N/A	N/A
2015/0106562	12/2014	Helmer	N/A	N/A
2015/0221349	12/2014	Jeon et al.	N/A	N/A
2015/0234709	12/2014	Koarashi	N/A	N/A
2015/0242856	12/2014	Dhurandhar et al.	N/A	N/A
2015/0244775	12/2014	Vibhor et al.	N/A	N/A
2015/0278534	12/2014	Thiyagarajan et al.	N/A	N/A
2015/0355957	12/2014	Steiner et al.	N/A	N/A
2015/0365427	12/2014	Ben-Shalom et al.	N/A	N/A
2015/0381646	12/2014	Lin	N/A	N/A
2016/0019114	12/2015	Han et al.	N/A	N/A
2016/0028754	12/2015	Mota et al.	N/A	N/A
2016/0078080	12/2015	Chen et al.	N/A	N/A
2016/0078245	12/2015	Amarendran et al.	N/A	N/A
2016/0098191	12/2015	Golden et al.	N/A	N/A
2016/0098199	12/2015	Golden et al.	N/A	N/A
2016/0099963	12/2015	Mahaffey et al.	N/A	N/A
2016/0112397	12/2015	Mankovskii	N/A	N/A
2016/0164918	12/2015	Satish et al.	N/A	N/A
2016/0170823	12/2015	Miller et al.	N/A	N/A
2016/0173599	12/2015	Chablani et al.	N/A	N/A
2016/0277423	12/2015	Apostolescu et al.	N/A	N/A
2016/0321455	12/2015	Deng	N/A	G06F 3/0685
2016/0323105	12/2015	Lee et al.	N/A	N/A
2016/0352518	12/2015	Ford et al.	N/A	N/A
2016/0352766	12/2015	Flacher et al.	N/A	N/A
2016/0357483	12/2015	Okita et al.	N/A	N/A
2016/0366156	12/2015	Kantor et al.	N/A	N/A
2016/0371152	12/2015	Parshin et al.	N/A	N/A
2016/0378352	12/2015	Khan et al.	N/A	N/A
2017/0034189	12/2016	Powell	N/A	N/A
2017/0054686	12/2016	Malkov et al.	N/A	N/A
2017/0061322	12/2016	Chari et al.	N/A	N/A
2017/0063534	12/2016	Brown	N/A	N/A
2017/0063887	12/2016	Muddu et al.	N/A	N/A

2017/0063888	12/2016	Muddu et al.	N/A	N/A
2017/0063907	12/2016	Muddu et al.	N/A	N/A
2017/0063908	12/2016	Muddu et al.	N/A	N/A
2017/0063909	12/2016	Muddu et al.	N/A	N/A
2017/0063910	12/2016	Muddu et al.	N/A	N/A
2017/0063911	12/2016	Muddu et al.	N/A	N/A
2017/0063912	12/2016	Muddu et al.	N/A	N/A
2017/0104773	12/2016	Flacher et al.	N/A	N/A
2017/0109057	12/2016	Karaca	N/A	N/A
2017/0118240	12/2016	Reddy et al.	N/A	N/A
2017/0131934	12/2016	Kaczmarczyk et al.	N/A	N/A
2017/0134403	12/2016	Hearn et al.	N/A	N/A
2017/0134410	12/2016	Muddu et al.	N/A	N/A
2017/0134415	12/2016	Muddu et al.	N/A	N/A
2017/0140156	12/2016	Gu et al.	N/A	N/A
2017/0142140	12/2016	Muddu et al.	N/A	N/A
2017/0147457	12/2016	Abhijeet et al.	N/A	N/A
2017/0149730	12/2016	Baughman et al.	N/A	N/A
2017/0155652	12/2016	Most et al.	N/A	N/A
2017/0177867	12/2016	Crofton et al.	N/A	N/A
2017/0180394	12/2016	Crofton et al.	N/A	N/A
2017/0185296	12/2016	Dong et al.	N/A	N/A
2017/0185301	12/2016	McGiverin et al.	N/A	N/A
2017/0192940	12/2016	Ghatage et al.	N/A	N/A
2017/0206353	12/2016	Jai et al.	N/A	N/A
2017/0213128	12/2016	Hammond et al.	N/A	N/A
2017/0213154	12/2016	Hammond et al.	N/A	N/A
2017/0213155	12/2016	Hammond et al.	N/A	N/A
2017/0213156	12/2016	Hammond et al.	N/A	N/A
2017/0223031	12/2016	Gu et al.	N/A	N/A
2017/0235497	12/2016	Shih	N/A	N/A
2017/0235507	12/2016	Sinha et al.	N/A	N/A
2017/0235561	12/2016	Butler et al.	N/A	N/A
2017/0235758	12/2016	Venkatesh et al.	N/A	N/A
2017/0242788	12/2016	Ahmed et al.	N/A	N/A
2017/0255652	12/2016	Kong et al.	N/A	N/A
2017/0256325	12/2016	Pappu et al.	N/A	N/A
2017/0277596	12/2016	Kyathanahalli et al.	N/A	N/A
2017/0279836	12/2016	Vasseur et al.	N/A	N/A
2017/0300422	12/2016	Szubocsev	N/A	N/A
2017/0301220	12/2016	Jarrell et al.	N/A	N/A
2017/0316210	12/2016	Patil et al.	N/A	N/A
2017/0324755	12/2016	Dekel et al.	N/A	N/A
2017/0329966	12/2016	Koganti et al.	N/A	N/A
2017/0331892	12/2016	Crofton et al.	N/A	N/A
2017/0346839	12/2016	Peppe et al.	N/A	N/A
2017/0364681	12/2016	Roguine et al.	N/A	N/A
2017/0374083	12/2016	Cohen et al.	N/A	N/A
2018/0004659	12/2017	Greenspan	N/A	N/A
2018/0007069	12/2017	Hunt et al.	N/A	N/A

2018/0024893	12/2017	Sella et al.	N/A	N/A
2018/0032758	12/2017	Wang et al.	N/A	N/A
2018/0048658	12/2017	Hittel et al.	N/A	N/A
2018/0052999	12/2017	Mitola	N/A	N/A
2018/0054452	12/2017	Muddu et al.	N/A	N/A
2018/0054454	12/2017	Astigarraga et al.	N/A	N/A
2018/0063179	12/2017	Salajegheh et al.	N/A	N/A
2018/0069888	12/2017	Muddu et al.	N/A	N/A
2018/0082075	12/2017	Copeland	N/A	N/A
2018/0103106	12/2017	Cocagne et al.	N/A	N/A
2018/0113638	12/2017	Petersen et al.	N/A	N/A
2018/0114234	12/2017	Fighel	N/A	N/A
2018/0115577	12/2017	Shukla et al.	N/A	N/A
2018/0124073	12/2017	Scherman et al.	N/A	N/A
2018/0124086	12/2017	Flacher et al.	N/A	N/A
2018/0129579	12/2017	Debnath et al.	N/A	N/A
2018/0139053	12/2017	Kadam et al.	N/A	N/A
2018/0157834	12/2017	Continella et al.	N/A	N/A
2018/0189000	12/2017	Li et al.	N/A	N/A
2018/0189508	12/2017	Li et al.	N/A	N/A
2018/0191763	12/2017	Hillard et al.	N/A	N/A
2018/0198765	12/2017	Maybee et al.	N/A	N/A
2018/0198812	12/2017	Christodorescu et al.	N/A	N/A
2018/0204000	12/2017	Charters et al.	N/A	N/A
2018/0211038	12/2017	Breiman et al.	N/A	N/A
2018/0211039	12/2017	Tamir et al.	N/A	N/A
2018/0211042	12/2017	Reinecke	N/A	G06F 21/568
2018/0212987	12/2017	Tamir et al.	N/A	N/A
2018/0217996	12/2017	Joshi et al.	N/A	N/A
2018/0225461	12/2017	Kotler et al.	N/A	N/A
2018/0239920	12/2017	Gupta	N/A	N/A
2018/0246649	12/2017	Datar et al.	N/A	N/A
2018/0248896	12/2017	Challita et al.	N/A	N/A
2018/0275902	12/2017	Monday et al.	N/A	N/A
2018/0278647	12/2017	Gabaev et al.	N/A	N/A
2018/0284758	12/2017	Cella et al.	N/A	N/A
2018/0285382	12/2017	Mehta et al.	N/A	N/A
2018/0285567	12/2017	Raman	N/A	N/A
2018/0293374	12/2017	Chen	N/A	N/A
2018/0295149	12/2017	Gazit et al.	N/A	N/A
2018/0307839	12/2017	Bhave et al.	N/A	N/A
2018/0351969	12/2017	MacLeod et al.	N/A	N/A
2018/0373877	12/2017	Bruso et al.	N/A	N/A
2018/0375826	12/2017	Chang et al.	N/A	N/A
2019/0007284	12/2018	Doshi et al.	N/A	N/A
2019/0012382	12/2018	Minkovicz et al.	N/A	N/A
2019/0018784	12/2018	Ishii et al.	N/A	N/A
2019/0034647	12/2018	Avrahami et al.	N/A	N/A
2019/0036958	12/2018	Shi	N/A	N/A

2019/0041842	12/2018	Cella et al.	N/A	N/A
2019/0042744	12/2018	Rajasekharan et al.	N/A	N/A
2019/0081876	12/2018	Ghare et al.	N/A	N/A
2019/0081981	12/2018	Bansal	N/A	N/A
2019/0087256	12/2018	Horrell et al.	N/A	N/A
2019/0087572	12/2018	Ellam et al.	N/A	N/A
2019/0098040	12/2018	Newman et al.	N/A	N/A
2019/0102262	12/2018	Sukhomlinov et al.	N/A	N/A
2019/0108099	12/2018	Mazumdar	N/A	N/A
2019/0108119	12/2018	Gholamipour et al.	N/A	N/A
2019/0108340	12/2018	Bedhapudi et al.	N/A	N/A
2019/0108341	12/2018	Bedhapudi et al.	N/A	N/A
2019/0108889	12/2018	Gholamipour et al.	N/A	N/A
2019/0109870	12/2018	Bedhapudi et al.	N/A	N/A
2019/0121576	12/2018	Jean	N/A	N/A
2019/0121972	12/2018	Norvill et al.	N/A	N/A
2019/0130097	12/2018	Berler et al.	N/A	N/A
2019/0138727	12/2018	Dontov et al.	N/A	N/A
2019/0156861	12/2018	Burton et al.	N/A	N/A
2019/0158512	12/2018	Zhang	N/A	N/A
2019/0188065	12/2018	Anghel et al.	N/A	N/A
2019/0188385	12/2018	Selvaraj et al.	N/A	N/A
2019/0196731	12/2018	Sapuntzakis et al.	N/A	N/A
2019/0197239	12/2018	Govardhan	N/A	N/A
2019/0207960	12/2018	Chu et al.	N/A	N/A
2019/0209022	12/2018	Sobol et al.	N/A	N/A
2019/0213137	12/2018	Lin et al.	N/A	N/A
2019/0258426	12/2018	Roh	N/A	N/A
2019/0266335	12/2018	Robinson	N/A	N/A
2019/0288840	12/2018	Gallancy et al.	N/A	N/A
2019/0294507	12/2018	Linnen et al.	N/A	N/A
2019/0303571	12/2018	Chelarescu et al.	N/A	N/A
2019/0303572	12/2018	Chelarescu et al.	N/A	N/A
2019/0303573	12/2018	Chelarescu et al.	N/A	N/A
2019/0303575	12/2018	Chelarescu et al.	N/A	N/A
2019/0319980	12/2018	Levy et al.	N/A	N/A
2019/0319987	12/2018	Levy et al.	N/A	N/A
2019/0324644	12/2018	Jin et al.	N/A	N/A
2019/0347418	12/2018	Strogov et al.	N/A	N/A
2019/0354443	12/2018	Haustein et al.	N/A	N/A
2019/0369918	12/2018	Cho et al.	N/A	N/A
2019/0392417	12/2018	Li et al.	N/A	N/A
2020/0019470	12/2019	Wolfson et al.	N/A	N/A
2020/0036736	12/2019	Novikov et al.	N/A	N/A
2020/0042707	12/2019	Kucherov et al.	N/A	N/A
2020/0065483	12/2019	Mu et al.	N/A	N/A
2020/0076831	12/2019	Baughman et al.	N/A	N/A
2020/0128027	12/2019	Hittel et al.	N/A	N/A
2020/0128028	12/2019	Hittel et al.	N/A	N/A
2020/0134049	12/2019	Bassov et al.	N/A	N/A

2020/0159674	12/2019	Morgan et al.	N/A	N/A
2020/0186602	12/2019	Leshinsky et al.	N/A	N/A
2020/0201777	12/2019	Kalamatianos et al.	N/A	N/A
2020/0204589	12/2019	Strogov et al.	N/A	N/A
2020/0233959	12/2019	Spurlock et al.	N/A	N/A
2020/0236085	12/2019	Spurlock et al.	N/A	N/A
2020/0236121	12/2019	Spurlock et al.	N/A	N/A
2020/0244678	12/2019	Shua	N/A	G06F 16/128
2020/0250051	12/2019	Chen et al.	N/A	N/A
2020/0250522	12/2019	Meiri et al.	N/A	N/A
2020/0257798	12/2019	Rao et al.	N/A	N/A
2020/0257820	12/2019	Li	N/A	N/A
2020/0274887	12/2019	Zaw	N/A	N/A
2020/0285737	12/2019	Kraus et al.	N/A	N/A
2020/0293696	12/2019	Motwani et al.	N/A	N/A
2020/0319979	12/2019	Kulaga et al.	N/A	N/A
2020/0341685	12/2019	Sharon et al.	N/A	N/A
2020/0344254	12/2019	Kokko et al.	N/A	N/A
2020/0356686	12/2019	Vijayvargiya et al.	N/A	N/A
2020/0358780	12/2019	Anbalagan	N/A	N/A
2020/0364128	12/2019	Vittal et al.	N/A	N/A
2020/0364429	12/2019	Yang	N/A	N/A
2020/0389469	12/2019	Litichever et al.	N/A	N/A
2021/0006575	12/2020	McLean et al.	N/A	N/A
2021/0011654	12/2020	Lee et al.	N/A	N/A
2021/0026961	12/2020	Underwood et al.	N/A	N/A
2021/0056203	12/2020	Qiao et al.	N/A	N/A
2021/0096957	12/2020	Rahman	N/A	G06F 11/1471
2021/0111881	12/2020	Migault et al.	N/A	N/A
2021/0117377	12/2020	Savir	N/A	G06F 16/178
2021/0133327	12/2020	Wu	N/A	G06F 21/577
2021/0133328	12/2020	Wu	N/A	G06F 8/65
2021/0152520	12/2020	Rigor et al.	N/A	N/A
2021/0152592	12/2020	Martynenko et al.	N/A	N/A
2021/0157523	12/2020	Matsushita et al.	N/A	N/A
2021/0182397	12/2020	Kamik et al.	N/A	N/A
2021/0216478	12/2020	Miller et al.	N/A	N/A
2021/0216625	12/2020	Miller et al.	N/A	N/A
2021/0216630	12/2020	Karr	N/A	N/A
2021/0216633	12/2020	Lee et al.	N/A	N/A
2021/0216646	12/2020	Miller et al.	N/A	N/A
2021/0216648	12/2020	Child et al.	N/A	N/A
2021/0216666	12/2020	Miller et al.	N/A	N/A
2021/0224379	12/2020	Pientka et al.	N/A	N/A
2021/0240608	12/2020	Jean et al.	N/A	N/A

2021/0255938	12/2020	Baker	N/A	G06F 11/1469
2021/0271757	12/2020	Horspool et al.	N/A	N/A
2021/0314341	12/2020	Moskovich et al.	N/A	N/A
2021/0319103	12/2020	Constantinescu et al.	N/A	N/A
2021/0374096	12/2020	Hein et al.	N/A	N/A
2021/0397711	12/2020	Karr	N/A	N/A
2022/0150241	12/2021	Nadiminti	N/A	H04L 63/0853
2022/0229909	12/2021	Shachar et al.	N/A	N/A
2022/0239690	12/2021	Arnon et al.	N/A	N/A
2024/0028748	12/2023	Marcon	N/A	G06F 3/0652

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
103370685	12/2012	CN	N/A
103370686	12/2015	CN	N/A
104025010	12/2015	CN	N/A
3066610	12/2015	EP	N/A
3082047	12/2015	EP	N/A
3120235	12/2016	EP	N/A
2007087036	12/2006	JP	N/A
2007094472	12/2006	JP	N/A
2008250667	12/2007	JP	N/A
2010211681	12/2009	JP	N/A
1995002349	12/1994	WO	N/A
1999013403	12/1998	WO	N/A
2008102347	12/2007	WO	N/A
2010071655	12/2009	WO	N/A

OTHER PUBLICATIONS

Hwang, et al., “RAID-x: A New Distributed Disk Array for I/O-Centric Cluster Computing”, Proceedings the Ninth International Symposium on High-Performance Distributed Computing, 2000, pp. 279-286. cited by applicant

International Search Report and Written Opinion received in International Application No. PCT/US2022/047039 dated Jan. 30, 2023. cited by applicant

Kharaz, A., et al., “Unveil: A Large-Scale, Automated Approach to Detecting Ransomware”, 25th USENIX Security Symposium, Aug. 10-12, 2016, Austin, TX. ISBN 978-1-931971-32-4. <https://www.usenix.org/conference/usenixsecurity16/technical-sessions/presentation/kharaz>. cited by applicant

Mehnaz, S., et al., “RWGuard: A Real-Time Detection System Against Cryptographic Ransomware”, Springer Nature Switzerland AG 2018. RAID LNCS 11050, pp. 114-136, 2018. https://doi.org/10.1007/978-3-030-00470-5_6. cited by applicant

Microsoft Corporation, “Fundamentals of Garbage Collection”, Retrieved Aug. 30, 2013 via the WayBack Machine, 11 pages. cited by applicant

Microsoft Corporation, “GCSettings.IsServerGC Property”, Retrieved Oct. 27, 2013 via the WayBack Machine, 3 pages. cited by applicant

Scaife, N., et al., “CryptoLock (and Drop It): Stopping Ransomware Attacks on User Data”, 2016

IEEE 36th International Conference on Distributed Computing Systems. cited by applicant
Stalzer, “FlashBlades: System Architecture and Applications”, Proceedings of the 2nd Workshop on Architectures and Systems for Big Data, Association for Computing Machinery, New York, NY, 2012, pp. 10-14. cited by applicant
Storer, et al., “Pergamum: Replacing Tape with Energy Efficient, Reliable, Disc-Based Archival Storage”, FAST’08: Proceedings of the 6th USENIX Conference on File and Storage Technologies, Feb. 2008, Article No. 1, pp. 1-16. cited by applicant
Strunk, J.D., et al., “Self-Securing Storage: Protecting Data in Compromised Systems”, Carnegie Mellon University. Operating Systems Design and Implementation (San Diego, CA, Oct. 23-25, 2000), pp. 165-180. USENIX Association, 2000. cited by applicant
International Search Report and Written Opinion for International Application No. PCT/US2023/023859, mailed Sep. 20, 2023, 13 pages. cited by applicant

Primary Examiner: Gracia; Gary S

Attorney, Agent or Firm: ALG Intellectual Property, LLC

Background/Summary

RELATED APPLICATIONS (1) This application is a continuation-in-part of U.S. patent application Ser. No. 18/127,926, filed Mar. 29, 2023, which is a continuation-in-part of U.S. patent application Ser. No. 17/039,486, filed Sep. 30, 2020 and issued as U.S. Pat. No. 11,720,692, U.S. patent application Ser. No. 17/039,536, filed Sep. 30, 2020 and issued as U.S. Pat. No. 11,625,481, U.S. patent application Ser. No. 17/039,556, filed Sep. 30, 2020 and issued as U.S. Pat. No. 11,720,714, U.S. patent application Ser. No. 17/039,604, filed Sep. 30, 2020 and issued as U.S. Pat. No. 11,651,075, U.S. patent application Ser. No. 17/074,313, filed Oct. 19, 2020, U.S. patent application Ser. No. 17/235,737, filed Apr. 20, 2021 and issued as U.S. Pat. No. 11,687,418, U.S. patent application Ser. No. 17/342,203, filed Jun. 8, 2021 and issued as U.S. Pat. No. 11,657,155, U.S. patent application Ser. No. 17/409,124, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/409,130, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/409,135, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/463,088, filed Aug. 31, 2021, U.S. patent application Ser. No. 17/506,501, filed Oct. 20, 2021, U.S. patent application Ser. No. 17/541,870, filed Dec. 3, 2021, U.S. patent application Ser. No. 17/506,509, filed Oct. 20, 2021, U.S. patent application Ser. No. 17/723,903, filed Apr. 19, 2022, U.S. patent application Ser. No. 17/725,182, filed Apr. 20, 2022 and issued as U.S. Pat. No. 11,657,146, U.S. patent application Ser. No. 17/846,301, filed Jun. 22, 2022, and to U.S. patent application Ser. No. 17/980,354, filed Nov. 3, 2022 and issued as U.S. Pat. No. 11,720,691, each of which is incorporated herein by reference in its entirety. (2) U.S. patent application Ser. No. 17/039,486 is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. U.S. patent application Ser. No. 17/039,486 also claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/985,229, filed Mar. 4, 2020, which is incorporated herein by reference in its entirety. (3) U.S. patent application Ser. No. 17/039,536 is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. U.S. patent application Ser. No. 17/039,536 also claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/985,229, filed Mar. 4, 2020, which is incorporated herein by reference in

[illegible]

Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. (14) U.S. patent application Ser. No. 17/541,870 is a continuation-in-part of U.S. patent application Ser. No. 16/916,903, filed Jun. 30, 2020, now U.S. Pat. No. 11,341,236, which application is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. (15) U.S. patent application Ser. No. 17/506,509 is a continuation-in-part of U.S. patent application Ser. No. 16/916,903, filed Jun. 30, 2020, now U.S. Pat. No. 11,341,236, which application is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. (16) U.S. patent application Ser. No. 17/723,903 is a continuation-in-part of U.S. patent application Ser. No. 16/916,903, filed Jun. 30, 2020, now U.S. Pat. No. 11,341,236, which application is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019. U.S. patent application Ser. No. 16/916,903 also claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/985,229, filed Mar. 4, 2020, each of which is incorporated herein by reference in its entirety. (17) U.S. patent application Ser. No. 17/725,182 is a continuation of U.S. patent application Ser. No. 16/916,903, filed Jun. 30, 2020, now U.S. Pat. No. 11,341,236, which is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. U.S. patent application Ser. No. 16/916,903 also claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/985,229, filed Mar. 4, 2020, each of which is incorporated herein by reference in its entirety. (18) U.S. patent application Ser. No. 17/846,301 is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed on Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. (19) U.S. patent application Ser. No. 17/980,354 is a continuation of U.S. patent application Ser. No. 17/161,553, filed Jan. 28, 2021, now U.S. Pat. No. 11,520,907, which is a continuation-in-part of U.S. patent application Ser. No. 16/917,030, filed Jun. 30, 2020 and issued as U.S. Pat. No. 11,675,898, which is a continuation-in-part of U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, which claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/939,518, filed Nov. 22, 2019, each of which is incorporated herein by reference in its entirety. U.S. patent application Ser. No. 17/161,553 also claims priority under 35 U.S.C. § 119(e) to U.S. Provisional Patent Application No. 62/985,229, filed Mar. 4, 2020, each of which is incorporated herein by reference in its entirety.

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The accompanying drawings illustrate various embodiments and are a part of the specification. The illustrated embodiments are merely examples and do not limit the scope of the disclosure. Throughout the drawings, identical or similar reference numbers designate identical or similar elements.

(2) FIG. 1A illustrates a first example system for data storage in accordance with some implementations.

(3) FIG. 1B illustrates a second example system for data storage in accordance with some implementations.

(4) FIG. 1C illustrates a third example system for data storage in accordance with some implementations.

(5) FIG. 1D illustrates a fourth example system for data storage in accordance with some

implementations.

- (6) FIG. 2A is a perspective view of a storage cluster with multiple storage nodes and internal storage coupled to each storage node to provide network attached storage, in accordance with some embodiments.
- (7) FIG. 2B is a block diagram showing an interconnect switch coupling multiple storage nodes in accordance with some embodiments.
- (8) FIG. 2C is a multiple level block diagram, showing contents of a storage node and contents of one of the non-volatile solid state storage units in accordance with some embodiments.
- (9) FIG. 2D shows a storage server environment, which uses embodiments of the storage nodes and storage units of some previous figures in accordance with some embodiments.
- (10) FIG. 2E is a blade hardware block diagram, showing a control plane, compute and storage planes, and authorities interacting with underlying physical resources, in accordance with some embodiments.
- (11) FIG. 2F depicts elasticity software layers in blades of a storage cluster, in accordance with some embodiments.
- (12) FIG. 2G depicts authorities and storage resources in blades of a storage cluster, in accordance with some embodiments.
- (13) FIG. 3A sets forth a diagram of a storage system that is coupled for data communications with a cloud services provider in accordance with some embodiments of the present disclosure.
- (14) FIG. 3B sets forth a diagram of a storage system in accordance with some embodiments of the present disclosure.
- (15) FIG. 3C sets forth an example of a cloud-based storage system in accordance with some embodiments of the present disclosure.
- (16) FIG. 3D illustrates an exemplary computing device that may be specifically configured to perform one or more of the processes described herein.
- (17) FIG. 4 illustrates an exemplary data protection system in accordance with some embodiments of the present disclosure.
- (18) FIG. 5 illustrates an exemplary configuration in which a storage system processes read traffic and write traffic in accordance with some embodiments of the present disclosure.
- (19) FIG. 6 shows an exemplary configuration in which a cloud-based monitoring system is communicatively coupled to storage system by way of a network in accordance with some embodiments of the present disclosure.
- (20) FIGS. 7-31 illustrate exemplary methods in accordance with some embodiments of the present disclosure.
- (21) FIG. 32 shows a method that may be performed by a data protection system.
- (22) FIG. 33 shows operations that may be performed by a data protection system as part of performing a remedial action with respect to a managed run-time environment.

Description

DESCRIPTION OF EMBODIMENTS

- (1) FIG. 1A illustrates an example system for data storage, in accordance with some implementations. System **100** (also referred to as “storage system” herein) includes numerous elements for purposes of illustration rather than limitation. It may be noted that system **100** may include the same, more, or fewer elements configured in the same or different manner in other implementations.
- (2) System **100** includes a number of computing devices **164A-B**. Computing devices (also referred to as “client devices” herein) may be embodied, for example, a server in a data center, a workstation, a personal computer, a notebook, or the like. Computing devices **164A-B** may be

coupled for data communications to one or more storage arrays **102A-B** through a storage area network ('SAN') **158** or a local area network ('LAN') **160**.

(3) The SAN **158** may be implemented with a variety of data communications fabrics, devices, and protocols. For example, the fabrics for SAN **158** may include Fibre Channel, Ethernet, Infiniband, Serial Attached Small Computer System Interface ('SAS'), or the like. Data communications protocols for use with SAN **158** may include Advanced Technology Attachment ('ATA'), Fibre Channel Protocol, Small Computer System Interface ('SCSI'), Internet Small Computer System Interface ('iSCSI'), HyperSCSI, Non-Volatile Memory Express ('NVMe') over Fabrics, or the like. It may be noted that SAN **158** is provided for illustration, rather than limitation. Other data communication couplings may be implemented between computing devices **164A-B** and storage arrays **102A-B**.

(4) The LAN **160** may also be implemented with a variety of fabrics, devices, and protocols. For example, the fabrics for LAN **160** may include Ethernet (802.3), wireless (802.11), or the like. Data communication protocols for use in LAN **160** may include Transmission Control Protocol ('TCP'), User Datagram Protocol ('UDP'), Internet Protocol ('IP'), HyperText Transfer Protocol ('HTTP'), Wireless Access Protocol ('WAP'), Handheld Device Transport Protocol ('HOTP'), Session Initiation Protocol ('SIP'), Real Time Protocol ('RTP'), or the like.

(5) Storage arrays **102A-B** may provide persistent data storage for the computing devices **164A-B**. Storage array **102A** may be contained in a chassis (not shown), and storage array **102B** may be contained in another chassis (not shown), in implementations. Storage array **102A** and **102B** may include one or more storage array controllers **110A-D** (also referred to as "controller" herein). A storage array controller **110A-D** may be embodied as a module of automated computing machinery comprising computer hardware, computer software, or a combination of computer hardware and software. In some implementations, the storage array controllers **110A-D** may be configured to carry out various storage tasks. Storage tasks may include writing data received from the computing devices **164A-B** to storage array **102A-B**, erasing data from storage array **102A-B**, retrieving data from storage array **102A-B** and providing data to computing devices **164A-B**, monitoring and reporting of disk utilization and performance, performing redundancy operations, such as Redundant Array of Independent Drives ('RAID') or RAID-like data redundancy operations, compressing data, encrypting data, and so forth.

(6) Storage array controller **110A-D** may be implemented in a variety of ways, including as a Field Programmable Gate Array ('FPGA'), a Programmable Logic Chip ('PLC'), an Application Specific Integrated Circuit ('ASIC'), System-on-Chip ('SOC'), or any computing device that includes discrete components such as a processing device, central processing unit, computer memory, or various adapters. Storage array controller **110A-D** may include, for example, a data communications adapter configured to support communications via the SAN **158** or LAN **160**. In some implementations, storage array controller **110A-D** may be independently coupled to the LAN **160**. In implementations, storage array controller **110A-D** may include an I/O controller or the like that couples the storage array controller **110A-D** for data communications, through a midplane (not shown), to a persistent storage resource **170A-B** (also referred to as a "storage resource" herein). The persistent storage resource **170A-B** may include any number of storage drives **171A-F** (also referred to as "storage devices" herein) and any number of non-volatile Random Access Memory ('NVRAM') devices (not shown).

(7) In some implementations, the NVRAM devices of a persistent storage resource **170A-B** may be configured to receive, from the storage array controller **110A-D**, data to be stored in the storage drives **171A-F**. In some examples, the data may originate from computing devices **164A-B**. In some examples, writing data to the NVRAM device may be carried out more quickly than directly writing data to the storage drive **171A-F**. In implementations, the storage array controller **110A-D** may be configured to utilize the NVRAM devices as a quickly accessible buffer for data destined to be written to the storage drives **171A-F**. Latency for write requests using NVRAM devices as a

buffer may be improved relative to a system in which a storage array controller **110A-D** writes data directly to the storage drives **171A-F**. In some implementations, the NVRAM devices may be implemented with computer memory in the form of high bandwidth, low latency RAM. The NVRAM device is referred to as “non-volatile” because the NVRAM device may receive or include a unique power source that maintains the state of the RAM after main power loss to the NVRAM device. Such a power source may be a battery, one or more capacitors, or the like. In response to a power loss, the NVRAM device may be configured to write the contents of the RAM to a persistent storage, such as the storage drives **171A-F**.

(8) In implementations, storage drive **171A-F** may refer to any device configured to record data persistently, where “persistently” or “persistent” refers as to a device's ability to maintain recorded data after loss of power. In some implementations, storage drive **171A-F** may correspond to non-disk storage media. For example, the storage drive **171A-F** may be one or more solid-state drives (‘SSDs’), flash memory based storage, any type of solid-state non-volatile memory, or any other type of non-mechanical storage device. In other implementations, storage drive **171A-F** may include mechanical or spinning hard disk, such as hard-disk drives (‘HDD’).

(9) In some implementations, the storage array controllers **110A-D** may be configured for offloading device management responsibilities from storage drive **171A-F** in storage array **102A-B**. For example, storage array controllers **110A-D** may manage control information that may describe the state of one or more memory blocks in the storage drives **171A-F**. The control information may indicate, for example, that a particular memory block has failed and should no longer be written to, that a particular memory block contains boot code for a storage array controller **110A-D**, the number of program-erase (‘P/E’) cycles that have been performed on a particular memory block, the age of data stored in a particular memory block, the type of data that is stored in a particular memory block, and so forth. In some implementations, the control information may be stored with an associated memory block as metadata. In other implementations, the control information for the storage drives **171A-F** may be stored in one or more particular memory blocks of the storage drives **171A-F** that are selected by the storage array controller **110A-D**. The selected memory blocks may be tagged with an identifier indicating that the selected memory block contains control information. The identifier may be utilized by the storage array controllers **110A-D** in conjunction with storage drives **171A-F** to quickly identify the memory blocks that contain control information. For example, the storage controllers **110A-D** may issue a command to locate memory blocks that contain control information. It may be noted that control information may be so large that parts of the control information may be stored in multiple locations, that the control information may be stored in multiple locations for purposes of redundancy, for example, or that the control information may otherwise be distributed across multiple memory blocks in the storage drive **171A-F**.

(10) In implementations, storage array controllers **110A-D** may offload device management responsibilities from storage drives **171A-F** of storage array **102A-B** by retrieving, from the storage drives **171A-F**, control information describing the state of one or more memory blocks in the storage drives **171A-F**. Retrieving the control information from the storage drives **171A-F** may be carried out, for example, by the storage array controller **110A-D** querying the storage drives **171A-F** for the location of control information for a particular storage drive **171A-F**. The storage drives **171A-F** may be configured to execute instructions that enable the storage drive **171A-F** to identify the location of the control information. The instructions may be executed by a controller (not shown) associated with or otherwise located on the storage drive **171A-F** and may cause the storage drive **171A-F** to scan a portion of each memory block to identify the memory blocks that store control information for the storage drives **171A-F**. The storage drives **171A-F** may respond by sending a response message to the storage array controller **110A-D** that includes the location of control information for the storage drive **171A-F**. Responsive to receiving the response message, storage array controllers **110A-D** may issue a request to read data stored at the address associated

with the location of control information for the storage drives **171A-F**.

(11) In other implementations, the storage array controllers **110A-D** may further offload device management responsibilities from storage drives **171A-F** by performing, in response to receiving the control information, a storage drive management operation. A storage drive management operation may include, for example, an operation that is typically performed by the storage drive **171A-F** (e.g., the controller (not shown) associated with a particular storage drive **171A-F**). A storage drive management operation may include, for example, ensuring that data is not written to failed memory blocks within the storage drive **171A-F**, ensuring that data is written to memory blocks within the storage drive **171A-F** in such a way that adequate wear leveling is achieved, and so forth.

(12) In implementations, storage array **102A-B** may implement two or more storage array controllers **110A-D**. For example, storage array **102A** may include storage array controllers **110A** and storage array controllers **110B**. At a given instance, a single storage array controller **110A-D** (e.g., storage array controller **110A**) of a storage system **100** may be designated with primary status (also referred to as “primary controller” herein), and other storage array controllers **110A-D** (e.g., storage array controller **110A**) may be designated with secondary status (also referred to as “secondary controller” herein). The primary controller may have particular rights, such as permission to alter data in persistent storage resource **170A-B** (e.g., writing data to persistent storage resource **170A-B**). At least some of the rights of the primary controller may supersede the rights of the secondary controller. For instance, the secondary controller may not have permission to alter data in persistent storage resource **170A-B** when the primary controller has the right. The status of storage array controllers **110A-D** may change. For example, storage array controller **110A** may be designated with secondary status, and storage array controller **110B** may be designated with primary status.

(13) In some implementations, a primary controller, such as storage array controller **110A**, may serve as the primary controller for one or more storage arrays **102A-B**, and a second controller, such as storage array controller **110B**, may serve as the secondary controller for the one or more storage arrays **102A-B**. For example, storage array controller **110A** may be the primary controller for storage array **102A** and storage array **102B**, and storage array controller **110B** may be the secondary controller for storage array **102A** and **102B**. In some implementations, storage array controllers **110C** and **110D** (also referred to as “storage processing modules”) may neither have primary or secondary status. Storage array controllers **110C** and **110D**, implemented as storage processing modules, may act as a communication interface between the primary and secondary controllers (e.g., storage array controllers **110A** and **110B**, respectively) and storage array **102B**. For example, storage array controller **110A** of storage array **102A** may send a write request, via SAN **158**, to storage array **102B**. The write request may be received by both storage array controllers **110C** and **110D** of storage array **102B**. Storage array controllers **110C** and **110D** facilitate the communication, e.g., send the write request to the appropriate storage drive **171A-F**. It may be noted that in some implementations storage processing modules may be used to increase the number of storage drives controlled by the primary and secondary controllers.

(14) In implementations, storage array controllers **110A-D** are communicatively coupled, via a midplane (not shown), to one or more storage drives **171A-F** and to one or more NVRAM devices (not shown) that are included as part of a storage array **102A-B**. The storage array controllers **110A-D** may be coupled to the midplane via one or more data communication links and the midplane may be coupled to the storage drives **171A-F** and the NVRAM devices via one or more data communications links. The data communications links described herein are collectively illustrated by data communications links **108A-D** and may include a Peripheral Component Interconnect Express (‘PCIe’) bus, for example.

(15) FIG. 1B illustrates an example system for data storage, in accordance with some implementations. Storage array controller **101** illustrated in FIG. 1B may be similar to the storage

array controllers **110A-D** described with respect to FIG. **1A**. In one example, storage array controller **101** may be similar to storage array controller **110A** or storage array controller **110B**. Storage array controller **101** includes numerous elements for purposes of illustration rather than limitation. It may be noted that storage array controller **101** may include the same, more, or fewer elements configured in the same or different manner in other implementations. It may be noted that elements of FIG. **1A** may be included below to help illustrate features of storage array controller **101**.

(16) Storage array controller **101** may include one or more processing devices **104** and random access memory ('RAM') **111**. Processing device **104** (or controller **101**) represents one or more general-purpose processing devices such as a microprocessor, central processing unit, or the like. More particularly, the processing device **104** (or controller **101**) may be a complex instruction set computing ('CISC') microprocessor, reduced instruction set computing ('RISC') microprocessor, very long instruction word ('VLIW') microprocessor, or a processor implementing other instruction sets or processors implementing a combination of instruction sets. The processing device **104** (or controller **101**) may also be one or more special-purpose processing devices such as an ASIC, an FPGA, a digital signal processor ('DSP'), network processor, or the like.

(17) The processing device **104** may be connected to the RAM **111** via a data communications link **106**, which may be embodied as a high speed memory bus such as a Double-Data Rate 4 ('DDR4') bus. Stored in RAM **111** is an operating system **112**. In some implementations, instructions **113** are stored in RAM **111**. Instructions **113** may include computer program instructions for performing operations in a direct-mapped flash storage system. In one embodiment, a direct-mapped flash storage system is one that addresses data blocks within flash drives directly and without an address translation performed by the storage controllers of the flash drives.

(18) In implementations, storage array controller **101** includes one or more host bus adapters **103A-C** that are coupled to the processing device **104** via a data communications link **105A-C**. In implementations, host bus adapters **103A-C** may be computer hardware that connects a host system (e.g., the storage array controller) to other network and storage arrays. In some examples, host bus adapters **103A-C** may be a Fibre Channel adapter that enables the storage array controller **101** to connect to a SAN, an Ethernet adapter that enables the storage array controller **101** to connect to a LAN, or the like. Host bus adapters **103A-C** may be coupled to the processing device **104** via a data communications link **105A-C** such as, for example, a PCIe bus.

(19) In implementations, storage array controller **101** may include a host bus adapter **114** that is coupled to an expander **115**. The expander **115** may be used to attach a host system to a larger number of storage drives. The expander **115** may, for example, be a SAS expander utilized to enable the host bus adapter **114** to attach to storage drives in an implementation where the host bus adapter **114** is embodied as a SAS controller.

(20) In implementations, storage array controller **101** may include a switch **116** coupled to the processing device **104** via a data communications link **109**. The switch **116** may be a computer hardware device that can create multiple endpoints out of a single endpoint, thereby enabling multiple devices to share a single endpoint. The switch **116** may, for example, be a PCIe switch that is coupled to a PCIe bus (e.g., data communications link **109**) and presents multiple PCIe connection points to the midplane.

(21) In implementations, storage array controller **101** includes a data communications link **107** for coupling the storage array controller **101** to other storage array controllers. In some examples, data communications link **107** may be a QuickPath Interconnect (QPI) interconnect.

(22) A traditional storage system that uses traditional flash drives may implement a process across the flash drives that are part of the traditional storage system. For example, a higher level process of the storage system may initiate and control a process across the flash drives. However, a flash drive of the traditional storage system may include its own storage controller that also performs the process. Thus, for the traditional storage system, a higher level process (e.g., initiated by the

storage system) and a lower level process (e.g., initiated by a storage controller of the storage system) may both be performed.

(23) To resolve various deficiencies of a traditional storage system, operations may be performed by higher level processes and not by the lower level processes. For example, the flash storage system may include flash drives that do not include storage controllers that provide the process. Thus, the operating system of the flash storage system itself may initiate and control the process. This may be accomplished by a direct-mapped flash storage system that addresses data blocks within the flash drives directly and without an address translation performed by the storage controllers of the flash drives.

(24) The operating system of the flash storage system may identify and maintain a list of allocation units across multiple flash drives of the flash storage system. The allocation units may be entire erase blocks or multiple erase blocks. The operating system may maintain a map or address range that directly maps addresses to erase blocks of the flash drives of the flash storage system.

(25) Direct mapping to the erase blocks of the flash drives may be used to rewrite data and erase data. For example, the operations may be performed on one or more allocation units that include a first data and a second data where the first data is to be retained and the second data is no longer being used by the flash storage system. The operating system may initiate the process to write the first data to new locations within other allocation units and erasing the second data and marking the allocation units as being available for use for subsequent data. Thus, the process may only be performed by the higher level operating system of the flash storage system without an additional lower level process being performed by controllers of the flash drives.

(26) Advantages of the process being performed only by the operating system of the flash storage system include increased reliability of the flash drives of the flash storage system as unnecessary or redundant write operations are not being performed during the process. One possible point of novelty here is the concept of initiating and controlling the process at the operating system of the flash storage system. In addition, the process can be controlled by the operating system across multiple flash drives. This is contrast to the process being performed by a storage controller of a flash drive.

(27) A storage system can consist of two storage array controllers that share a set of drives for failover purposes, or it could consist of a single storage array controller that provides a storage service that utilizes multiple drives, or it could consist of a distributed network of storage array controllers each with some number of drives or some amount of Flash storage where the storage array controllers in the network collaborate to provide a complete storage service and collaborate on various aspects of a storage service including storage allocation and garbage collection.

(28) FIG. 1C illustrates a third example system **117** for data storage in accordance with some implementations. System **117** (also referred to as “storage system” herein) includes numerous elements for purposes of illustration rather than limitation. It may be noted that system **117** may include the same, more, or fewer elements configured in the same or different manner in other implementations.

(29) In one embodiment, system **117** includes a dual Peripheral Component Interconnect (‘PCI’) flash storage device **118** with separately addressable fast write storage. System **117** may include a storage controller **119**. In one embodiment, storage controller **119A-D** may be a CPU, ASIC, FPGA, or any other circuitry that may implement control structures necessary according to the present disclosure. In one embodiment, system **117** includes flash memory devices (e.g., including flash memory devices **120a-n**), operatively coupled to various channels of the storage device controller **119**. Flash memory devices **120a-n**, may be presented to the controller **119A-D** as an addressable collection of Flash pages, erase blocks, and/or control elements sufficient to allow the storage device controller **119A-D** to program and retrieve various aspects of the Flash. In one embodiment, storage device controller **119A-D** may perform operations on flash memory devices **120a-n** including storing and retrieving data content of pages, arranging and erasing any blocks,

tracking statistics related to the use and reuse of Flash memory pages, erase blocks, and cells, tracking and predicting error codes and faults within the Flash memory, controlling voltage levels associated with programming and retrieving contents of Flash cells, etc.

(30) In one embodiment, system **117** may include RAM **121** to store separately addressable fast-write data. In one embodiment, RAM **121** may be one or more separate discrete devices. In another embodiment, RAM **121** may be integrated into storage device controller **119A-D** or multiple storage device controllers. The RAM **121** may be utilized for other purposes as well, such as temporary program memory for a processing device (e.g., a CPU) in the storage device controller **119**.

(31) In one embodiment, system **117** may include a stored energy device **122**, such as a rechargeable battery or a capacitor. Stored energy device **122** may store energy sufficient to power the storage device controller **119**, some amount of the RAM (e.g., RAM **121**), and some amount of Flash memory (e.g., Flash memory **120a-120n**) for sufficient time to write the contents of RAM to Flash memory. In one embodiment, storage device controller **119A-D** may write the contents of RAM to Flash Memory if the storage device controller detects loss of external power.

(32) In one embodiment, system **117** includes two data communications links **123a**, **123b**. In one embodiment, data communications links **123a**, **123b** may be PCI interfaces. In another embodiment, data communications links **123a**, **123b** may be based on other communications standards (e.g., HyperTransport, InfiniBand, etc.). Data communications links **123a**, **123b** may be based on non-volatile memory express ('NVMe') or NVMe over fabrics ('NVMe over fabrics') specifications that allow external connection to the storage device controller **119A-D** from other components in the storage system **117**. It should be noted that data communications links may be interchangeably referred to herein as PCI buses for convenience.

(33) System **117** may also include an external power source (not shown), which may be provided over one or both data communications links **123a**, **123b**, or which may be provided separately. An alternative embodiment includes a separate Flash memory (not shown) dedicated for use in storing the content of RAM **121**. The storage device controller **119A-D** may present a logical device over a PCI bus which may include an addressable fast-write logical device, or a distinct part of the logical address space of the storage device **118**, which may be presented as PCI memory or as persistent storage. In one embodiment, operations to store into the device are directed into the RAM **121**. On power failure, the storage device controller **119A-D** may write stored content associated with the addressable fast-write logical storage to Flash memory (e.g., Flash memory **120a-n**) for long-term persistent storage.

(34) In one embodiment, the logical device may include some presentation of some or all of the content of the Flash memory devices **120a-n**, where that presentation allows a storage system including a storage device **118** (e.g., storage system **117**) to directly address Flash memory pages and directly reprogram erase blocks from storage system components that are external to the storage device through the PCI bus. The presentation may also allow one or more of the external components to control and retrieve other aspects of the Flash memory including some or all of: tracking statistics related to use and reuse of Flash memory pages, erase blocks, and cells across all the Flash memory devices; tracking and predicting error codes and faults within and across the Flash memory devices; controlling voltage levels associated with programming and retrieving contents of Flash cells; etc.

(35) In one embodiment, the stored energy device **122** may be sufficient to ensure completion of in-progress operations to the Flash memory devices **120a-120n** stored energy device **122** may power storage device controller **119A-D** and associated Flash memory devices (e.g., **120a-n**) for those operations, as well as for the storing of fast-write RAM to Flash memory. Stored energy device **122** may be used to store accumulated statistics and other parameters kept and tracked by the Flash memory devices **120a-n** and/or the storage device controller **119**. Separate capacitors or stored energy devices (such as smaller capacitors near or embedded within the Flash memory devices

themselves) may be used for some or all of the operations described herein.

(36) Various schemes may be used to track and optimize the life span of the stored energy component, such as adjusting voltage levels over time, partially discharging the storage energy device **122** to measure corresponding discharge characteristics, etc. If the available energy decreases over time, the effective available capacity of the addressable fast-write storage may be decreased to ensure that it can be written safely based on the currently available stored energy.

(37) FIG. **1D** illustrates a third example system **124** for data storage in accordance with some implementations. In one embodiment, system **124** includes storage controllers **125a**, **125b**. In one embodiment, storage controllers **125a**, **125b** are operatively coupled to Dual PCI storage devices **119a**, **119b** and **119c**, **119d**, respectively. Storage controllers **125a**, **125b** may be operatively coupled (e.g., via a storage network **130**) to some number of host computers **127a-n**.

(38) In one embodiment, two storage controllers (e.g., **125a** and **125b**) provide storage services, such as a SCS) block storage array, a file server, an object server, a database or data analytics service, etc. The storage controllers **125a**, **125b** may provide services through some number of network interfaces (e.g., **126a-d**) to host computers **127a-n** outside of the storage system **124**. Storage controllers **125a**, **125b** may provide integrated services or an application entirely within the storage system **124**, forming a converged storage and compute system. The storage controllers **125a**, **125b** may utilize the fast write memory within or across storage devices **119a-d** to journal in progress operations to ensure the operations are not lost on a power failure, storage controller removal, storage controller or storage system shutdown, or some fault of one or more software or hardware components within the storage system **124**.

(39) In one embodiment, controllers **125a**, **125b** operate as PCI masters to one or the other PCI buses **128a**, **128b**. In another embodiment, **128a** and **128b** may be based on other communications standards (e.g., HyperTransport, InfiniBand, etc.). Other storage system embodiments may operate storage controllers **125a**, **125b** as multi-masters for both PCI buses **128a**, **128b**. Alternately, a PCI/NVMe/NVMf switching infrastructure or fabric may connect multiple storage controllers. Some storage system embodiments may allow storage devices to communicate with each other directly rather than communicating only with storage controllers. In one embodiment, a storage device controller **119a** may be operable under direction from a storage controller **125a** to synthesize and transfer data to be stored into Flash memory devices from data that has been stored in RAM (e.g., RAM **121** of FIG. **1C**). For example, a recalculated version of RAM content may be transferred after a storage controller has determined that an operation has fully committed across the storage system, or when fast-write memory on the device has reached a certain used capacity, or after a certain amount of time, to ensure improve safety of the data or to release addressable fast-write capacity for reuse. This mechanism may be used, for example, to avoid a second transfer over a bus (e.g., **128a**, **128b**) from the storage controllers **125a**, **125b**. In one embodiment, a recalculation may include compressing data, attaching indexing or other metadata, combining multiple data segments together, performing erasure code calculations, etc.

(40) In one embodiment, under direction from a storage controller **125a**, **125b**, a storage device controller **119a**, **119b** may be operable to calculate and transfer data to other storage devices from data stored in RAM (e.g., RAM **121** of FIG. **1C**) without involvement of the storage controllers **125a**, **125b**. This operation may be used to mirror data stored in one controller **125a** to another controller **125b**, or it could be used to offload compression, data aggregation, and/or erasure coding calculations and transfers to storage devices to reduce load on storage controllers or the storage controller interface **129a**, **129b** to the PCI bus **128a**, **128b**.

(41) A storage device controller **119A-D** may include mechanisms for implementing high availability primitives for use by other parts of a storage system external to the Dual PCI storage device **118**. For example, reservation or exclusion primitives may be provided so that, in a storage system with two storage controllers providing a highly available storage service, one storage controller may prevent the other storage controller from accessing or continuing to access the

storage device. This could be used, for example, in cases where one controller detects that the other controller is not functioning properly or where the interconnect between the two storage controllers may itself not be functioning properly.

(42) In one embodiment, a storage system for use with Dual PCI direct mapped storage devices with separately addressable fast write storage includes systems that manage erase blocks or groups of erase blocks as allocation units for storing data on behalf of the storage service, or for storing metadata (e.g., indexes, logs, etc.) associated with the storage service, or for proper management of the storage system itself. Flash pages, which may be a few kilobytes in size, may be written as data arrives or as the storage system is to persist data for long intervals of time (e.g., above a defined threshold of time). To commit data more quickly, or to reduce the number of writes to the Flash memory devices, the storage controllers may first write data into the separately addressable fast write storage on one more storage devices.

(43) In one embodiment, the storage controllers **125a**, **125b** may initiate the use of erase blocks within and across storage devices (e.g., **118**) in accordance with an age and expected remaining lifespan of the storage devices, or based on other statistics. The storage controllers **125a**, **125b** may initiate garbage collection and data migration data between storage devices in accordance with pages that are no longer needed as well as to manage Flash page and erase block lifespans and to manage overall system performance.

(44) In one embodiment, the storage system **124** may utilize mirroring and/or erasure coding schemes as part of storing data into addressable fast write storage and/or as part of writing data into allocation units associated with erase blocks. Erasure codes may be used across storage devices, as well as within erase blocks or allocation units, or within and across Flash memory devices on a single storage device, to provide redundancy against single or multiple storage device failures or to protect against internal corruptions of Flash memory pages resulting from Flash memory operations or from degradation of Flash memory cells. Mirroring and erasure coding at various levels may be used to recover from multiple types of failures that occur separately or in combination.

(45) The embodiments depicted with reference to FIGS. 2A-G illustrate a storage cluster that stores user data, such as user data originating from one or more user or client systems or other sources external to the storage cluster. The storage cluster distributes user data across storage nodes housed within a chassis, or across multiple chassis, using erasure coding and redundant copies of metadata. Erasure coding refers to a method of data protection or reconstruction in which data is stored across a set of different locations, such as disks, storage nodes or geographic locations. Flash memory is one type of solid-state memory that may be integrated with the embodiments, although the embodiments may be extended to other types of solid-state memory or other storage medium, including non-solid state memory. Control of storage locations and workloads are distributed across the storage locations in a clustered peer-to-peer system. Tasks such as mediating communications between the various storage nodes, detecting when a storage node has become unavailable, and balancing I/Os (inputs and outputs) across the various storage nodes, are all handled on a distributed basis. Data is laid out or distributed across multiple storage nodes in data fragments or stripes that support data recovery in some embodiments. Ownership of data can be reassigned within a cluster, independent of input and output patterns. This architecture described in more detail below allows a storage node in the cluster to fail, with the system remaining operational, since the data can be reconstructed from other storage nodes and thus remain available for input and output operations. In various embodiments, a storage node may be referred to as a cluster node, a blade, or a server.

(46) The storage cluster may be contained within a chassis, i.e., an enclosure housing one or more storage nodes. A mechanism to provide power to each storage node, such as a power distribution bus, and a communication mechanism, such as a communication bus that enables communication between the storage nodes are included within the chassis. The storage cluster can run as an

independent system in one location according to some embodiments. In one embodiment, a chassis contains at least two instances of both the power distribution and the communication bus which may be enabled or disabled independently. The internal communication bus may be an Ethernet bus, however, other technologies such as PCIe, InfiniBand, and others, are equally suitable. The chassis provides a port for an external communication bus for enabling communication between multiple chassis, directly or through a switch, and with client systems. The external communication may use a technology such as Ethernet, InfiniBand, Fibre Channel, etc. In some embodiments, the external communication bus uses different communication bus technologies for inter-chassis and client communication. If a switch is deployed within or between chassis, the switch may act as a translation between multiple protocols or technologies. When multiple chassis are connected to define a storage cluster, the storage cluster may be accessed by a client using either proprietary interfaces or standard interfaces such as network file system ('NFS'), common internet file system ('CIFS'), small computer system interface ('SCSI') or hypertext transfer protocol ('HTTP'). Translation from the client protocol may occur at the switch, chassis external communication bus or within each storage node. In some embodiments, multiple chassis may be coupled or connected to each other through an aggregator switch. A portion and/or all of the coupled or connected chassis may be designated as a storage cluster. As discussed above, each chassis can have multiple blades, each blade has a media access control ('MAC') address, but the storage cluster is presented to an external network as having a single cluster IP address and a single MAC address in some embodiments.

(47) Each storage node may be one or more storage servers and each storage server is connected to one or more non-volatile solid state memory units, which may be referred to as storage units or storage devices. One embodiment includes a single storage server in each storage node and between one to eight non-volatile solid state memory units, however this one example is not meant to be limiting. The storage server may include a processor, DRAM and interfaces for the internal communication bus and power distribution for each of the power buses. Inside the storage node, the interfaces and storage unit share a communication bus, e.g., PCI Express, in some embodiments. The non-volatile solid state memory units may directly access the internal communication bus interface through a storage node communication bus, or request the storage node to access the bus interface. The non-volatile solid state memory unit contains an embedded CPU, solid state storage controller, and a quantity of solid state mass storage, e.g., between 2-32 terabytes ('TB') in some embodiments. An embedded volatile storage medium, such as DRAM, and an energy reserve apparatus are included in the non-volatile solid state memory unit. In some embodiments, the energy reserve apparatus is a capacitor, super-capacitor, or battery that enables transferring a subset of DRAM contents to a stable storage medium in the case of power loss. In some embodiments, the non-volatile solid state memory unit is constructed with a storage class memory, such as phase change or magnetoresistive random access memory ('MRAM') that substitutes for DRAM and enables a reduced power hold-up apparatus.

(48) One of many features of the storage nodes and non-volatile solid state storage is the ability to proactively rebuild data in a storage cluster. The storage nodes and non-volatile solid state storage can determine when a storage node or non-volatile solid state storage in the storage cluster is unreachable, independent of whether there is an attempt to read data involving that storage node or non-volatile solid state storage. The storage nodes and non-volatile solid state storage then cooperate to recover and rebuild the data in at least partially new locations. This constitutes a proactive rebuild, in that the system rebuilds data without waiting until the data is needed for a read access initiated from a client system employing the storage cluster. These and further details of the storage memory and operation thereof are discussed below.

(49) FIG. 2A is a perspective view of a storage cluster **161**, with multiple storage nodes **150** and internal solid-state memory coupled to each storage node to provide network attached storage or storage area network, in accordance with some embodiments. A network attached storage, storage

area network, or a storage cluster, or other storage memory, could include one or more storage clusters **161**, each having one or more storage nodes **150**, in a flexible and reconfigurable arrangement of both the physical components and the amount of storage memory provided thereby. The storage cluster **161** is designed to fit in a rack, and one or more racks can be set up and populated as desired for the storage memory. The storage cluster **161** has a chassis **138** having multiple slots **142**. It should be appreciated that chassis **138** may be referred to as a housing, enclosure, or rack unit. In one embodiment, the chassis **138** has fourteen slots **142**, although other numbers of slots are readily devised. For example, some embodiments have four slots, eight slots, sixteen slots, thirty-two slots, or other suitable number of slots. Each slot **142** can accommodate one storage node **150** in some embodiments. Chassis **138** includes flaps **148** that can be utilized to mount the chassis **138** on a rack. Fans **144** provide air circulation for cooling of the storage nodes **150** and components thereof, although other cooling components could be used, or an embodiment could be devised without cooling components. A switch fabric **146** couples storage nodes **150** within chassis **138** together and to a network for communication to the memory. In an embodiment depicted in herein, the slots **142** to the left of the switch fabric **146** and fans **144** are shown occupied by storage nodes **150**, while the slots **142** to the right of the switch fabric **146** and fans **144** are empty and available for insertion of storage node **150** for illustrative purposes. This configuration is one example, and one or more storage nodes **150** could occupy the slots **142** in various further arrangements. The storage node arrangements need not be sequential or adjacent in some embodiments. Storage nodes **150** are hot pluggable, meaning that a storage node **150** can be inserted into a slot **142** in the chassis **138**, or removed from a slot **142**, without stopping or powering down the system. Upon insertion or removal of storage node **150** from slot **142**, the system automatically reconfigures in order to recognize and adapt to the change. Reconfiguration, in some embodiments, includes restoring redundancy and/or rebalancing data or load.

(50) Each storage node **150** can have multiple components. In the embodiment shown here, the storage node **150** includes a printed circuit board **159** populated by a CPU **156**, i.e., processor, a memory **154** coupled to the CPU **156**, and a non-volatile solid state storage **152** coupled to the CPU **156**, although other mountings and/or components could be used in further embodiments. The memory **154** has instructions which are executed by the CPU **156** and/or data operated on by the CPU **156**. As further explained below, the non-volatile solid state storage **152** includes flash or, in further embodiments, other types of solid-state memory.

(51) Referring to FIG. 2A, storage cluster **161** is scalable, meaning that storage capacity with non-uniform storage sizes is readily added, as described above. One or more storage nodes **150** can be plugged into or removed from each chassis and the storage cluster self-configures in some embodiments. Plug-in storage nodes **150**, whether installed in a chassis as delivered or later added, can have different sizes. For example, in one embodiment a storage node **150** can have any multiple of 4 TB, e.g., 8 TB, 12 TB, 16 TB, 32 TB, etc. In further embodiments, a storage node **150** could have any multiple of other storage amounts or capacities. Storage capacity of each storage node **150** is broadcast, and influences decisions of how to stripe the data. For maximum storage efficiency, an embodiment can self-configure as wide as possible in the stripe, subject to a predetermined requirement of continued operation with loss of up to one, or up to two, non-volatile solid state storage units **152** or storage nodes **150** within the chassis.

(52) FIG. 2B is a block diagram showing a communications interconnect **173** and power distribution bus **172** coupling multiple storage nodes **150**. Referring back to FIG. 2A, the communications interconnect **173** can be included in or implemented with the switch fabric **146** in some embodiments. Where multiple storage clusters **161** occupy a rack, the communications interconnect **173** can be included in or implemented with a top of rack switch, in some embodiments. As illustrated in FIG. 2B, storage cluster **161** is enclosed within a single chassis **138**. External port **176** is coupled to storage nodes **150** through communications interconnect **173**, while external port **174** is coupled directly to a storage node. External power port **178** is coupled to power

distribution bus **172**. Storage nodes **150** may include varying amounts and differing capacities of non-volatile solid state storage **152** as described with reference to FIG. 2A. In addition, one or more storage nodes **150** may be a compute only storage node as illustrated in FIG. 2B. Authorities **168** are implemented on the non-volatile solid state storages **152**, for example as lists or other data structures stored in memory. In some embodiments the authorities are stored within the non-volatile solid state storage **152** and supported by software executing on a controller or other processor of the non-volatile solid state storage **152**. In a further embodiment, authorities **168** are implemented on the storage nodes **150**, for example as lists or other data structures stored in the memory **154** and supported by software executing on the CPU **156** of the storage node **150**. Authorities **168** control how and where data is stored in the non-volatile solid state storages **152** in some embodiments. This control assists in determining which type of erasure coding scheme is applied to the data, and which storage nodes **150** have which portions of the data. Each authority **168** may be assigned to a non-volatile solid state storage **152**. Each authority may control a range of inode numbers, segment numbers, or other data identifiers which are assigned to data by a file system, by the storage nodes **150**, or by the non-volatile solid state storage **152**, in various embodiments.

(53) Every piece of data, and every piece of metadata, has redundancy in the system in some embodiments. In addition, every piece of data and every piece of metadata has an owner, which may be referred to as an authority. If that authority is unreachable, for example through failure of a storage node, there is a plan of succession for how to find that data or that metadata. In various embodiments, there are redundant copies of authorities **168**. Authorities **168** have a relationship to storage nodes **150** and non-volatile solid state storage **152** in some embodiments. Each authority **168**, covering a range of data segment numbers or other identifiers of the data, may be assigned to a specific non-volatile solid state storage **152**. In some embodiments the authorities **168** for all of such ranges are distributed over the non-volatile solid state storages **152** of a storage cluster. Each storage node **150** has a network port that provides access to the non-volatile solid state storage(s) **152** of that storage node **150**. Data can be stored in a segment, which is associated with a segment number and that segment number is an indirection for a configuration of a RAID (redundant array of independent disks) stripe in some embodiments. The assignment and use of the authorities **168** thus establishes an indirection to data. Indirection may be referred to as the ability to reference data indirectly, in this case via an authority **168**, in accordance with some embodiments. A segment identifies a set of non-volatile solid state storage **152** and a local identifier into the set of non-volatile solid state storage **152** that may contain data. In some embodiments, the local identifier is an offset into the device and may be reused sequentially by multiple segments. In other embodiments the local identifier is unique for a specific segment and never reused. The offsets in the non-volatile solid state storage **152** are applied to locating data for writing to or reading from the non-volatile solid state storage **152** (in the form of a RAID stripe). Data is striped across multiple units of non-volatile solid state storage **152**, which may include or be different from the non-volatile solid state storage **152** having the authority **168** for a particular data segment.

(54) If there is a change in where a particular segment of data is located, e.g., during a data move or a data reconstruction, the authority **168** for that data segment should be consulted, at that non-volatile solid state storage **152** or storage node **150** having that authority **168**. In order to locate a particular piece of data, embodiments calculate a hash value for a data segment or apply an inode number or a data segment number. The output of this operation points to a non-volatile solid state storage **152** having the authority **168** for that particular piece of data. In some embodiments there are two stages to this operation. The first stage maps an entity identifier (ID), e.g., a segment number, inode number, or directory number to an authority identifier. This mapping may include a calculation such as a hash or a bit mask. The second stage is mapping the authority identifier to a particular non-volatile solid state storage **152**, which may be done through an explicit mapping. The operation is repeatable, so that when the calculation is performed, the result of the calculation repeatably and reliably points to a particular non-volatile solid state storage **152** having that

authority **168**. The operation may include the set of reachable storage nodes as input. If the set of reachable non-volatile solid state storage units changes the optimal set changes. In some embodiments, the persisted value is the current assignment (which is always true) and the calculated value is the target assignment the cluster will attempt to reconfigure towards. This calculation may be used to determine the optimal non-volatile solid state storage **152** for an authority in the presence of a set of non-volatile solid state storage **152** that are reachable and constitute the same cluster. The calculation also determines an ordered set of peer non-volatile solid state storage **152** that will also record the authority to non-volatile solid state storage mapping so that the authority may be determined even if the assigned non-volatile solid state storage is unreachable. A duplicate or substitute authority **168** may be consulted if a specific authority **168** is unavailable in some embodiments.

(55) With reference to FIGS. 2A and 2B, two of the many tasks of the CPU **156** on a storage node **150** are to break up write data, and reassemble read data. When the system has determined that data is to be written, the authority **168** for that data is located as above. When the segment ID for data is already determined the request to write is forwarded to the non-volatile solid state storage **152** currently determined to be the host of the authority **168** determined from the segment. The host CPU **156** of the storage node **150**, on which the non-volatile solid state storage **152** and corresponding authority **168** reside, then breaks up or shards the data and transmits the data out to various non-volatile solid state storage **152**. The transmitted data is written as a data stripe in accordance with an erasure coding scheme. In some embodiments, data is requested to be pulled, and in other embodiments, data is pushed. In reverse, when data is read, the authority **168** for the segment ID containing the data is located as described above. The host CPU **156** of the storage node **150** on which the non-volatile solid state storage **152** and corresponding authority **168** reside requests the data from the non-volatile solid state storage and corresponding storage nodes pointed to by the authority. In some embodiments the data is read from flash storage as a data stripe. The host CPU **156** of storage node **150** then reassembles the read data, correcting any errors (if present) according to the appropriate erasure coding scheme, and forwards the reassembled data to the network. In further embodiments, some or all of these tasks can be handled in the non-volatile solid state storage **152**. In some embodiments, the segment host requests the data be sent to storage node **150** by requesting pages from storage and then sending the data to the storage node making the original request.

(56) In some systems, for example in UNIX-style file systems, data is handled with an index node or inode, which specifies a data structure that represents an object in a file system. The object could be a file or a directory, for example. Metadata may accompany the object, as attributes such as permission data and a creation timestamp, among other attributes. A segment number could be assigned to all or a portion of such an object in a file system. In other systems, data segments are handled with a segment number assigned elsewhere. For purposes of discussion, the unit of distribution is an entity, and an entity can be a file, a directory or a segment. That is, entities are units of data or metadata stored by a storage system. Entities are grouped into sets called authorities. Each authority has an authority owner, which is a storage node that has the exclusive right to update the entities in the authority. In other words, a storage node contains the authority, and that the authority, in turn, contains entities.

(57) A segment is a logical container of data in accordance with some embodiments. A segment is an address space between medium address space and physical flash locations, i.e., the data segment number, are in this address space. Segments may also contain meta-data, which enable data redundancy to be restored (rewritten to different flash locations or devices) without the involvement of higher level software. In one embodiment, an internal format of a segment contains client data and medium mappings to determine the position of that data. Each data segment is protected, e.g., from memory and other failures, by breaking the segment into a number of data and parity shards, where applicable. The data and parity shards are distributed, i.e., striped, across non-

volatile solid state storage **152** coupled to the host CPUs **156** (See FIGS. 2E and 2G) in accordance with an erasure coding scheme. Usage of the term segments refers to the container and its place in the address space of segments in some embodiments. Usage of the term stripe refers to the same set of shards as a segment and includes how the shards are distributed along with redundancy or parity information in accordance with some embodiments.

(58) A series of address-space transformations takes place across an entire storage system. At the top are the directory entries (file names) which link to an inode. Inodes point into medium address space, where data is logically stored. Medium addresses may be mapped through a series of indirect mediums to spread the load of large files, or implement data services like deduplication or snapshots. Medium addresses may be mapped through a series of indirect mediums to spread the load of large files, or implement data services like deduplication or snapshots. Segment addresses are then translated into physical flash locations. Physical flash locations have an address range bounded by the amount of flash in the system in accordance with some embodiments. Medium addresses and segment addresses are logical containers, and in some embodiments use a 128 bit or larger identifier so as to be practically infinite, with a likelihood of reuse calculated as longer than the expected life of the system. Addresses from logical containers are allocated in a hierarchical fashion in some embodiments. Initially, each non-volatile solid state storage unit **152** may be assigned a range of address space. Within this assigned range, the non-volatile solid state storage **152** is able to allocate addresses without synchronization with other non-volatile solid state storage **152**.

(59) Data and metadata is stored by a set of underlying storage layouts that are optimized for varying workload patterns and storage devices. These layouts incorporate multiple redundancy schemes, compression formats and index algorithms. Some of these layouts store information about authorities and authority masters, while others store file metadata and file data. The redundancy schemes include error correction codes that tolerate corrupted bits within a single storage device (such as a NAND flash chip), erasure codes that tolerate the failure of multiple storage nodes, and replication schemes that tolerate data center or regional failures. In some embodiments, low density parity check ('LDPC') code is used within a single storage unit. Reed-Solomon encoding is used within a storage cluster, and mirroring is used within a storage grid in some embodiments. Metadata may be stored using an ordered log structured index (such as a Log Structured Merge Tree), and large data may not be stored in a log structured layout.

(60) In order to maintain consistency across multiple copies of an entity, the storage nodes agree implicitly on two things through calculations: (1) the authority that contains the entity, and (2) the storage node that contains the authority. The assignment of entities to authorities can be done by pseudo randomly assigning entities to authorities, by splitting entities into ranges based upon an externally produced key, or by placing a single entity into each authority. Examples of pseudorandom schemes are linear hashing and the Replication Under Scalable Hashing ('RUSH') family of hashes, including Controlled Replication Under Scalable Hashing ('CRUSH'). In some embodiments, pseudo-random assignment is utilized only for assigning authorities to nodes because the set of nodes can change. The set of authorities cannot change so any subjective function may be applied in these embodiments. Some placement schemes automatically place authorities on storage nodes, while other placement schemes rely on an explicit mapping of authorities to storage nodes. In some embodiments, a pseudorandom scheme is utilized to map from each authority to a set of candidate authority owners. A pseudorandom data distribution function related to CRUSH may assign authorities to storage nodes and create a list of where the authorities are assigned. Each storage node has a copy of the pseudorandom data distribution function, and can arrive at the same calculation for distributing, and later finding or locating an authority. Each of the pseudorandom schemes requires the reachable set of storage nodes as input in some embodiments in order to conclude the same target nodes. Once an entity has been placed in an authority, the entity may be stored on physical devices so that no expected failure will lead to

unexpected data loss. In some embodiments, rebalancing algorithms attempt to store the copies of all entities within an authority in the same layout and on the same set of machines.

(61) Examples of expected failures include device failures, stolen machines, datacenter fires, and regional disasters, such as nuclear or geological events. Different failures lead to different levels of acceptable data loss. In some embodiments, a stolen storage node impacts neither the security nor the reliability of the system, while depending on system configuration, a regional event could lead to no loss of data, a few seconds or minutes of lost updates, or even complete data loss.

(62) In the embodiments, the placement of data for storage redundancy is independent of the placement of authorities for data consistency. In some embodiments, storage nodes that contain authorities do not contain any persistent storage. Instead, the storage nodes are connected to non-volatile solid state storage units that do not contain authorities. The communications interconnect between storage nodes and non-volatile solid state storage units consists of multiple communication technologies and has non-uniform performance and fault tolerance characteristics. In some embodiments, as mentioned above, non-volatile solid state storage units are connected to storage nodes via PCI express, storage nodes are connected together within a single chassis using Ethernet backplane, and chassis are connected together to form a storage cluster. Storage clusters are connected to clients using Ethernet or fiber channel in some embodiments. If multiple storage clusters are configured into a storage grid, the multiple storage clusters are connected using the Internet or other long-distance networking links, such as a “metro scale” link or private link that does not traverse the internet.

(63) Authority owners have the exclusive right to modify entities, to migrate entities from one non-volatile solid state storage unit to another non-volatile solid state storage unit, and to add and remove copies of entities. This allows for maintaining the redundancy of the underlying data. When an authority owner fails, is going to be decommissioned, or is overloaded, the authority is transferred to a new storage node. Transient failures make it non-trivial to ensure that all non-faulty machines agree upon the new authority location. The ambiguity that arises due to transient failures can be achieved automatically by a consensus protocol such as Paxos, hot-warm failover schemes, via manual intervention by a remote system administrator, or by a local hardware administrator (such as by physically removing the failed machine from the cluster, or pressing a button on the failed machine). In some embodiments, a consensus protocol is used, and failover is automatic. If too many failures or replication events occur in too short a time period, the system goes into a self-preservation mode and halts replication and data movement activities until an administrator intervenes in accordance with some embodiments.

(64) As authorities are transferred between storage nodes and authority owners update entities in their authorities, the system transfers messages between the storage nodes and non-volatile solid state storage units. With regard to persistent messages, messages that have different purposes are of different types. Depending on the type of the message, the system maintains different ordering and durability guarantees. As the persistent messages are being processed, the messages are temporarily stored in multiple durable and non-durable storage hardware technologies. In some embodiments, messages are stored in RAM, NVRAM and on NAND flash devices, and a variety of protocols are used in order to make efficient use of each storage medium. Latency-sensitive client requests may be persisted in replicated NVRAM, and then later NAND, while background rebalancing operations are persisted directly to NAND.

(65) Persistent messages are persistently stored prior to being transmitted. This allows the system to continue to serve client requests despite failures and component replacement. Although many hardware components contain unique identifiers that are visible to system administrators, manufacturer, hardware supply chain and ongoing monitoring quality control infrastructure, applications running on top of the infrastructure address virtualize addresses. These virtualized addresses do not change over the lifetime of the storage system, regardless of component failures and replacements. This allows each component of the storage system to be replaced over time

without reconfiguration or disruptions of client request processing, i.e., the system supports non-disruptive upgrades.

(66) In some embodiments, the virtualized addresses are stored with sufficient redundancy. A continuous monitoring system correlates hardware and software status and the hardware identifiers. This allows detection and prediction of failures due to faulty components and manufacturing details. The monitoring system also enables the proactive transfer of authorities and entities away from impacted devices before failure occurs by removing the component from the critical path in some embodiments.

(67) FIG. 2C is a multiple level block diagram, showing contents of a storage node **150** and contents of a non-volatile solid state storage **152** of the storage node **150**. Data is communicated to and from the storage node **150** by a network interface controller ('NIC') **202** in some embodiments. Each storage node **150** has a CPU **156**, and one or more non-volatile solid state storage **152**, as discussed above. Moving down one level in FIG. 2C, each non-volatile solid state storage **152** has a relatively fast non-volatile solid state memory, such as nonvolatile random access memory ('NVRAM') **204**, and flash memory **206**. In some embodiments, NVRAM **204** may be a component that does not require program/erase cycles (DRAM, MRAM, PCM), and can be a memory that can support being written vastly more often than the memory is read from. Moving down another level in FIG. 2C, the NVRAM **204** is implemented in one embodiment as high speed volatile memory, such as dynamic random access memory (DRAM) **216**, backed up by energy reserve **218**. Energy reserve **218** provides sufficient electrical power to keep the DRAM **216** powered long enough for contents to be transferred to the flash memory **206** in the event of power failure. In some embodiments, energy reserve **218** is a capacitor, super-capacitor, battery, or other device, that supplies a suitable supply of energy sufficient to enable the transfer of the contents of DRAM **216** to a stable storage medium in the case of power loss. The flash memory **206** is implemented as multiple flash dies **222**, which may be referred to as packages of flash dies **222** or an array of flash dies **222**. It should be appreciated that the flash dies **222** could be packaged in any number of ways, with a single die per package, multiple dies per package (i.e. multichip packages), in hybrid packages, as bare dies on a printed circuit board or other substrate, as encapsulated dies, etc. In the embodiment shown, the non-volatile solid state storage **152** has a controller **212** or other processor, and an input output (I/O) port **210** coupled to the controller **212**. I/O port **210** is coupled to the CPU **156** and/or the network interface controller **202** of the flash storage node **150**. Flash input output (I/O) port **220** is coupled to the flash dies **222**, and a direct memory access unit (DMA) **214** is coupled to the controller **212**, the DRAM **216** and the flash dies **222**. In the embodiment shown, the I/O port **210**, controller **212**, DMA unit **214** and flash I/O port **220** are implemented on a programmable logic device ('PLD') **208**, e.g., an FPGA. In this embodiment, each flash die **222** has pages, organized as sixteen kB (kilobyte) pages **224**, and a register **226** through which data can be written to or read from the flash die **222**. In further embodiments, other types of solid-state memory are used in place of, or in addition to flash memory illustrated within flash die **222**.

(68) Storage clusters **161**, in various embodiments as disclosed herein, can be contrasted with storage arrays in general. The storage nodes **150** are part of a collection that creates the storage cluster **161**. Each storage node **150** owns a slice of data and computing required to provide the data. Multiple storage nodes **150** cooperate to store and retrieve the data. Storage memory or storage devices, as used in storage arrays in general, are less involved with processing and manipulating the data. Storage memory or storage devices in a storage array receive commands to read, write, or erase data. The storage memory or storage devices in a storage array are not aware of a larger system in which they are embedded, or what the data means. Storage memory or storage devices in storage arrays can include various types of storage memory, such as RAM, solid state drives, hard disk drives, etc. The storage units **152** described herein have multiple interfaces active simultaneously and serving multiple purposes. In some embodiments, some of the functionality of

a storage node **150** is shifted into a storage unit **152**, transforming the storage unit **152** into a combination of storage unit **152** and storage node **150**. Placing computing (relative to storage data) into the storage unit **152** places this computing closer to the data itself. The various system embodiments have a hierarchy of storage node layers with different capabilities. By contrast, in a storage array, a controller owns and knows everything about all of the data that the controller manages in a shelf or storage devices. In a storage cluster **161**, as described herein, multiple controllers in multiple storage units **152** and/or storage nodes **150** cooperate in various ways (e.g., for erasure coding, data sharding, metadata communication and redundancy, storage capacity expansion or contraction, data recovery, and so on).

(69) FIG. 2D shows a storage server environment, which uses embodiments of the storage nodes **150** and storage units **152** of FIGS. 2A-C. In this version, each storage unit **152** has a processor such as controller **212** (see FIG. 2C), an FPGA, flash memory **206**, and NVRAM **204** (which is super-capacitor backed DRAM **216**, see FIGS. 2B and 2C) on a PCIe (peripheral component interconnect express) board in a chassis **138** (see FIG. 2A). The storage unit **152** may be implemented as a single board containing storage, and may be the largest tolerable failure domain inside the chassis. In some embodiments, up to two storage units **152** may fail and the device will continue with no data loss.

(70) The physical storage is divided into named regions based on application usage in some embodiments. The NVRAM **204** is a contiguous block of reserved memory in the storage unit **152** DRAM **216**, and is backed by NAND flash. NVRAM **204** is logically divided into multiple memory regions written for two as spool (e.g., spool_region). Space within the NVRAM **204** spools is managed by each authority **168** independently. Each device provides an amount of storage space to each authority **168**. That authority **168** further manages lifetimes and allocations within that space. Examples of a spool include distributed transactions or notions. When the primary power to a storage unit **152** fails, onboard super-capacitors provide a short duration of power hold up. During this holdup interval, the contents of the NVRAM **204** are flushed to flash memory **206**. On the next power-on, the contents of the NVRAM **204** are recovered from the flash memory **206**.

(71) As for the storage unit controller, the responsibility of the logical “controller” is distributed across each of the blades containing authorities **168**. This distribution of logical control is shown in FIG. 2D as a host controller **242**, mid-tier controller **244** and storage unit controller(s) **246**. Management of the control plane and the storage plane are treated independently, although parts may be physically co-located on the same blade. Each authority **168** effectively serves as an independent controller. Each authority **168** provides its own data and metadata structures, its own background workers, and maintains its own lifecycle.

(72) FIG. 2E is a blade **252** hardware block diagram, showing a control plane **254**, compute and storage planes **256**, **258**, and authorities **168** interacting with underlying physical resources, using embodiments of the storage nodes **150** and storage units **152** of FIGS. 2A-C in the storage server environment of FIG. 2D. The control plane **254** is partitioned into a number of authorities **168** which can use the compute resources in the compute plane **256** to run on any of the blades **252**. The storage plane **258** is partitioned into a set of devices, each of which provides access to flash **206** and NVRAM **204** resources. In one embodiment, the compute plane **256** may perform the operations of a storage array controller, as described herein, on one or more devices of the storage plane **258** (e.g., a storage array).

(73) In the compute and storage planes **256**, **258** of FIG. 2E, the authorities **168** interact with the underlying physical resources (i.e., devices). From the point of view of an authority **168**, its resources are striped over all of the physical devices. From the point of view of a device, it provides resources to all authorities **168**, irrespective of where the authorities happen to run. Each authority **168** has allocated or has been allocated one or more partitions **260** of storage memory in the storage units **152**, e.g. partitions **260** in flash memory **206** and NVRAM **204**. Each authority **168** uses those allocated partitions **260** that belong to it, for writing or reading user data.

Authorities can be associated with differing amounts of physical storage of the system. For example, one authority **168** could have a larger number of partitions **260** or larger sized partitions **260** in one or more storage units **152** than one or more other authorities **168**.

(74) FIG. 2F depicts elasticity software layers in blades **252** of a storage cluster, in accordance with some embodiments. In the elasticity structure, elasticity software is symmetric, i.e., each blade's compute module **270** runs the three identical layers of processes depicted in FIG. 2F. Storage managers **274** execute read and write requests from other blades **252** for data and metadata stored in local storage unit **152** NVRAM **204** and flash **206**. Authorities **168** fulfill client requests by issuing the necessary reads and writes to the blades **252** on whose storage units **152** the corresponding data or metadata resides. Endpoints **272** parse client connection requests received from switch fabric **146** supervisory software, relay the client connection requests to the authorities **168** responsible for fulfillment, and relay the authorities' **168** responses to clients. The symmetric three-layer structure enables the storage system's high degree of concurrency. Elasticity scales out efficiently and reliably in these embodiments. In addition, elasticity implements a unique scale-out technique that balances work evenly across all resources regardless of client access pattern, and maximizes concurrency by eliminating much of the need for inter-blade coordination that typically occurs with conventional distributed locking.

(75) Still referring to FIG. 2F, authorities **168** running in the compute modules **270** of a blade **252** perform the internal operations required to fulfill client requests. One feature of elasticity is that authorities **168** are stateless, i.e., they cache active data and metadata in their own blades' **252** DRAMs for fast access, but the authorities store every update in their NVRAM **204** partitions on three separate blades **252** until the update has been written to flash **206**. All the storage system writes to NVRAM **204** are in triplicate to partitions on three separate blades **252** in some embodiments. With triple-mirrored NVRAM **204** and persistent storage protected by parity and Reed-Solomon RAID checksums, the storage system can survive concurrent failure of two blades **252** with no loss of data, metadata, or access to either.

(76) Because authorities **168** are stateless, they can migrate between blades **252**. Each authority **168** has a unique identifier. NVRAM **204** and flash **206** partitions are associated with authorities' **168** identifiers, not with the blades **252** on which they are running in some. Thus, when an authority **168** migrates, the authority **168** continues to manage the same storage partitions from its new location. When a new blade **252** is installed in an embodiment of the storage cluster, the system automatically rebalances load by: partitioning the new blade's **252** storage for use by the system's authorities **168**, migrating selected authorities **168** to the new blade **252**, starting endpoints **272** on the new blade **252** and including them in the switch fabric's **146** client connection distribution algorithm.

(77) From their new locations, migrated authorities **168** persist the contents of their NVRAM **204** partitions on flash **206**, process read and write requests from other authorities **168**, and fulfill the client requests that endpoints **272** direct to them. Similarly, if a blade **252** fails or is removed, the system redistributes its authorities **168** among the system's remaining blades **252**. The redistributed authorities **168** continue to perform their original functions from their new locations.

(78) FIG. 2G depicts authorities **168** and storage resources in blades **252** of a storage cluster, in accordance with some embodiments. Each authority **168** is exclusively responsible for a partition of the flash **206** and NVRAM **204** on each blade **252**. The authority **168** manages the content and integrity of its partitions independently of other authorities **168**. Authorities **168** compress incoming data and preserve it temporarily in their NVRAM **204** partitions, and then consolidate, RAID-protect, and persist the data in segments of the storage in their flash **206** partitions. As the authorities **168** write data to flash **206**, storage managers **274** perform the necessary flash translation to optimize write performance and maximize media longevity. In the background, authorities **168** “garbage collect,” or reclaim space occupied by data that clients have made obsolete by overwriting the data. It should be appreciated that since authorities' **168** partitions are

disjoint, there is no need for distributed locking to execute client and writes or to perform background functions.

(79) The embodiments described herein may utilize various software, communication and/or networking protocols. In addition, the configuration of the hardware and/or software may be adjusted to accommodate various protocols. For example, the embodiments may utilize Active Directory, which is a database based system that provides authentication, directory, policy, and other services in a WINDOWS™ environment. In these embodiments, LDAP (Lightweight Directory Access Protocol) is one example application protocol for querying and modifying items in directory service providers such as Active Directory. In some embodiments, a network lock manager ('NLM') is utilized as a facility that works in cooperation with the Network File System ('NFS') to provide a System V style of advisory file and record locking over a network. The Server Message Block ('SMB') protocol, one version of which is also known as Common Internet File System ('CIFS'), may be integrated with the storage systems discussed herein. SMP operates as an application-layer network protocol typically used for providing shared access to files, printers, and serial ports and miscellaneous communications between nodes on a network. SMB also provides an authenticated inter-process communication mechanism. AMAZON™ S3 (Simple Storage Service) is a web service offered by Amazon Web Services, and the systems described herein may interface with Amazon S3 through web services interfaces (REST (representational state transfer), SOAP (simple object access protocol), and BitTorrent). A RESTful API (application programming interface) breaks down a transaction to create a series of small modules. Each module addresses a particular underlying part of the transaction. The control or permissions provided with these embodiments, especially for object data, may include utilization of an access control list ('ACL'). The ACL is a list of permissions attached to an object and the ACL specifies which users or system processes are granted access to objects, as well as what operations are allowed on given objects. The systems may utilize Internet Protocol version 6 ('IPv6'), as well as IPv4, for the communications protocol that provides an identification and location system for computers on networks and routes traffic across the Internet. The routing of packets between networked systems may include Equal-cost multi-path routing ('ECMP'), which is a routing strategy where next-hop packet forwarding to a single destination can occur over multiple "best paths" which tie for top place in routing metric calculations. Multi-path routing can be used in conjunction with most routing protocols, because it is a per-hop decision limited to a single router. The software may support Multi-tenancy, which is an architecture in which a single instance of a software application serves multiple customers. Each customer may be referred to as a tenant. Tenants may be given the ability to customize some parts of the application, but may not customize the application's code, in some embodiments. The embodiments may maintain audit logs. An audit log is a document that records an event in a computing system. In addition to documenting what resources were accessed, audit log entries typically include destination and source addresses, a timestamp, and user login information for compliance with various regulations. The embodiments may support various key management policies, such as encryption key rotation. In addition, the system may support dynamic root passwords or some variation dynamically changing passwords.

(80) FIG. 3A sets forth a diagram of a storage system 306 that is coupled for data communications with a cloud services provider 302 in accordance with some embodiments of the present disclosure. Although depicted in less detail, the storage system 306 depicted in FIG. 3A may be similar to the storage systems described above with reference to FIGS. 1A-ID and FIGS. 2A-2G. In some embodiments, the storage system 306 depicted in FIG. 3A may be embodied as a storage system that includes imbalanced active/active controllers, as a storage system that includes balanced active/active controllers, as a storage system that includes active/active controllers where less than all of each controller's resources are utilized such that each controller has reserve resources that may be used to support failover, as a storage system that includes fully active/active controllers, as a storage system that includes dataset-segregated controllers, as a storage system that includes dual-

layer architectures with front-end controllers and back-end integrated storage controllers, as a storage system that includes scale-out clusters of dual-controller arrays, as well as combinations of such embodiments.

(81) In the example depicted in FIG. 3A, the storage system **306** is coupled to the cloud services provider **302** via a data communications link **304**. The data communications link **304** may be embodied as a dedicated data communications link, as a data communications pathway that is provided through the use of one or data communications networks such as a wide area network ('WAN') or LAN, or as some other mechanism capable of transporting digital information between the storage system **306** and the cloud services provider **302**. Such a data communications link **304** may be fully wired, fully wireless, or some aggregation of wired and wireless data communications pathways. In such an example, digital information may be exchanged between the storage system **306** and the cloud services provider **302** via the data communications link **304** using one or more data communications protocols. For example, digital information may be exchanged between the storage system **306** and the cloud services provider **302** via the data communications link **304** using the handheld device transfer protocol ('HDTP'), hypertext transfer protocol ('HTTP'), internet protocol ('IP'), real-time transfer protocol ('RTP'), transmission control protocol ('TCP'), user datagram protocol ('UDP'), wireless application protocol ('WAP'), or other protocol.

(82) The cloud services provider **302** depicted in FIG. 3A may be embodied, for example, as a system and computing environment that provides a vast array of services to users of the cloud services provider **302** through the sharing of computing resources via the data communications link **304**. The cloud services provider **302** may provide on-demand access to a shared pool of configurable computing resources such as computer networks, servers, storage, applications and services, and so on. The shared pool of configurable resources may be rapidly provisioned and released to a user of the cloud services provider **302** with minimal management effort. Generally, the user of the cloud services provider **302** is unaware of the exact computing resources utilized by the cloud services provider **302** to provide the services. Although in many cases such a cloud services provider **302** may be accessible via the Internet, readers of skill in the art will recognize that any system that abstracts the use of shared resources to provide services to a user through any data communications link may be considered a cloud services provider **302**.

(83) In the example depicted in FIG. 3A, the cloud services provider **302** may be configured to provide a variety of services to the storage system **306** and users of the storage system **306** through the implementation of various service models. For example, the cloud services provider **302** may be configured to provide services through the implementation of an infrastructure as a service ('IaaS') service model, through the implementation of a platform as a service ('PaaS') service model, through the implementation of a software as a service ('SaaS') service model, through the implementation of an authentication as a service ('AaaS') service model, through the implementation of a storage as a service model where the cloud services provider **302** offers access to its storage infrastructure for use by the storage system **306** and users of the storage system **306**, and so on. Readers will appreciate that the cloud services provider **302** may be configured to provide additional services to the storage system **306** and users of the storage system **306** through the implementation of additional service models, as the service models described above are included only for explanatory purposes and in no way represent a limitation of the services that may be offered by the cloud services provider **302** or a limitation as to the service models that may be implemented by the cloud services provider **302**.

(84) In the example depicted in FIG. 3A, the cloud services provider **302** may be embodied, for example, as a private cloud, as a public cloud, or as a combination of a private cloud and public cloud. In an embodiment in which the cloud services provider **302** is embodied as a private cloud, the cloud services provider **302** may be dedicated to providing services to a single organization rather than providing services to multiple organizations. In an embodiment where the cloud services provider **302** is embodied as a public cloud, the cloud services provider **302** may provide

services to multiple organizations. In still alternative embodiments, the cloud services provider **302** may be embodied as a mix of a private and public cloud services with a hybrid cloud deployment. (85) Although not explicitly depicted in FIG. 3A, readers will appreciate that a vast amount of additional hardware components and additional software components may be necessary to facilitate the delivery of cloud services to the storage system **306** and users of the storage system **306**. For example, the storage system **306** may be coupled to (or even include) a cloud storage gateway. Such a cloud storage gateway may be embodied, for example, as hardware-based or software-based appliance that is located on premise with the storage system **306**. Such a cloud storage gateway may operate as a bridge between local applications that are executing on the storage array **306** and remote, cloud-based storage that is utilized by the storage array **306**. Through the use of a cloud storage gateway, organizations may move primary iSCSI or NAS to the cloud services provider **302**, thereby enabling the organization to save space on their on-premises storage systems. Such a cloud storage gateway may be configured to emulate a disk array, a block-based device, a file server, or other storage system that can translate the SCSI commands, file server commands, or other appropriate command into REST-space protocols that facilitate communications with the cloud services provider **302**.

(86) In order to enable the storage system **306** and users of the storage system **306** to make use of the services provided by the cloud services provider **302**, a cloud migration process may take place during which data, applications, or other elements from an organization's local systems (or even from another cloud environment) are moved to the cloud services provider **302**. In order to successfully migrate data, applications, or other elements to the cloud services provider's **302** environment, middleware such as a cloud migration tool may be utilized to bridge gaps between the cloud services provider's **302** environment and an organization's environment. Such cloud migration tools may also be configured to address potentially high network costs and long transfer times associated with migrating large volumes of data to the cloud services provider **302**, as well as addressing security concerns associated with sensitive data to the cloud services provider **302** over data communications networks. In order to further enable the storage system **306** and users of the storage system **306** to make use of the services provided by the cloud services provider **302**, a cloud orchestrator may also be used to arrange and coordinate automated tasks in pursuit of creating a consolidated process or workflow. Such a cloud orchestrator may perform tasks such as configuring various components, whether those components are cloud components or on-premises components, as well as managing the interconnections between such components. The cloud orchestrator can simplify the inter-component communication and connections to ensure that links are correctly configured and maintained.

(87) In the example depicted in FIG. 3A, and as described briefly above, the cloud services provider **302** may be configured to provide services to the storage system **306** and users of the storage system **306** through the usage of a SaaS service model, eliminating the need to install and run the application on local computers, which may simplify maintenance and support of the application. Such applications may take many forms in accordance with various embodiments of the present disclosure. For example, the cloud services provider **302** may be configured to provide access to data analytics applications to the storage system **306** and users of the storage system **306**. Such data analytics applications may be configured, for example, to receive vast amounts of telemetry data phoned home by the storage system **306**. Such telemetry data may describe various operating characteristics of the storage system **306** and may be analyzed for a vast array of purposes including, for example, to determine the health of the storage system **306**, to identify workloads that are executing on the storage system **306**, to predict when the storage system **306** will run out of various resources, to recommend configuration changes, hardware or software upgrades, workflow migrations, or other actions that may improve the operation of the storage system **306**.

(88) The cloud services provider **302** may also be configured to provide access to virtualized

computing environments to the storage system **306** and users of the storage system **306**. Such virtualized computing environments may be embodied, for example, as a virtual machine or other virtualized computer hardware platforms, virtual storage devices, virtualized computer network resources, and so on. Examples of such virtualized environments can include virtual machines that are created to emulate an actual computer, virtualized desktop environments that separate a logical desktop from a physical machine, virtualized file systems that allow uniform access to different types of concrete file systems, and many others.

(89) For further explanation, FIG. **3B** sets forth a diagram of a storage system **306** in accordance with some embodiments of the present disclosure. Although depicted in less detail, the storage system **306** depicted in FIG. **3B** may be similar to the storage systems described above with reference to FIGS. **1A-1D** and FIGS. **2A-2G** as the storage system may include many of the components described above.

(90) The storage system **306** depicted in FIG. **3B** may include a vast amount of storage resources **308**, which may be embodied in many forms. For example, the storage resources **308** can include nano-RAM or another form of nonvolatile random access memory that utilizes carbon nanotubes deposited on a substrate, 3D crosspoint non-volatile memory, flash memory including single-level cell ('SLC') NAND flash, multi-level cell ('MLC') NAND flash, triple-level cell ('TLC') NAND flash, quad-level cell ('QLC') NAND flash, or others. Likewise, the storage resources **308** may include non-volatile magnetoresistive random-access memory ('MRAM'), including spin transfer torque ('STT') MRAM. The example storage resources **308** may alternatively include non-volatile phase-change memory ('PCM'), quantum memory that allows for the storage and retrieval of photonic quantum information, resistive random-access memory ('ReRAM'), storage class memory ('SCM'), or other form of storage resources, including any combination of resources described herein. Readers will appreciate that other forms of computer memories and storage devices may be utilized by the storage systems described above, including DRAM, SRAM, EEPROM, universal memory, and many others. The storage resources **308** depicted in FIG. **3A** may be embodied in a variety of form factors, including but not limited to, dual in-line memory modules ('DIMMs'), non-volatile dual in-line memory modules ('NVDIMMs'), M.2, U.2, and others.

(91) The storage resources **308** depicted in FIG. **3A** may include various forms of SCM. SCM may effectively treat fast, non-volatile memory (e.g., NAND flash) as an extension of DRAM such that an entire dataset may be treated as an in-memory dataset that resides entirely in DRAM. SCM may include non-volatile media such as, for example, NAND flash. Such NAND flash may be accessed utilizing NVMe that can use the PCIe bus as its transport, providing for relatively low access latencies compared to older protocols. In fact, the network protocols used for SSDs in all-flash arrays can include NVMe using Ethernet (ROCE, NVME TCP), Fibre Channel (NVMe FC), InfiniBand (iWARP), and others that make it possible to treat fast, non-volatile memory as an extension of DRAM. In view of the fact that DRAM is often byte-addressable and fast, non-volatile memory such as NAND flash is block-addressable, a controller software/hardware stack may be needed to convert the block data to the bytes that are stored in the media. Examples of media and software that may be used as SCM can include, for example, 3D XPoint, Intel Memory Drive Technology, Samsung's Z-SSD, and others.

(92) The example storage system **306** depicted in FIG. **3B** may implement a variety of storage architectures. For example, storage systems in accordance with some embodiments of the present disclosure may utilize block storage where data is stored in blocks, and each block essentially acts as an individual hard drive. Storage systems in accordance with some embodiments of the present disclosure may utilize object storage, where data is managed as objects. Each object may include the data itself, a variable amount of metadata, and a globally unique identifier, where object storage can be implemented at multiple levels (e.g., device level, system level, interface level). Storage systems in accordance with some embodiments of the present disclosure utilize file storage in

which data is stored in a hierarchical structure. Such data may be saved in files and folders, and presented to both the system storing it and the system retrieving it in the same format.

(93) The example storage system **306** depicted in FIG. **3B** may be embodied as a storage system in which additional storage resources can be added through the use of a scale-up model, additional storage resources can be added through the use of a scale-out model, or through some combination thereof. In a scale-up model, additional storage may be added by adding additional storage devices. In a scale-out model, however, additional storage nodes may be added to a cluster of storage nodes, where such storage nodes can include additional processing resources, additional networking resources, and so on.

(94) The storage system **306** depicted in FIG. **3B** also includes communications resources **310** that may be useful in facilitating data communications between components within the storage system **306**, as well as data communications between the storage system **306** and computing devices that are outside of the storage system **306**, including embodiments where those resources are separated by a relatively vast expanse. The communications resources **310** may be configured to utilize a variety of different protocols and data communication fabrics to facilitate data communications between components within the storage systems as well as computing devices that are outside of the storage system. For example, the communications resources **310** can include fibre channel ('FC') technologies such as FC fabrics and FC protocols that can transport SCSI commands over FC network, FC over ethernet ('FCoE') technologies through which FC frames are encapsulated and transmitted over Ethernet networks, InfiniBand ('B') technologies in which a switched fabric topology is utilized to facilitate transmissions between channel adapters, NVM Express ('NVMe') technologies and NVMe over fabrics ('NVMeoF') technologies through which non-volatile storage media attached via a PCI express ('PCIe') bus may be accessed, and others. In fact, the storage systems described above may, directly or indirectly, make use of neutrino communication technologies and devices through which information (including binary information) is transmitted using a beam of neutrinos.

(95) The communications resources **310** can also include mechanisms for accessing storage resources **308** within the storage system **306** utilizing serial attached SCSI ('SAS'), serial ATA ('SATA') bus interfaces for connecting storage resources **308** within the storage system **306** to host bus adapters within the storage system **306**, internet small computer systems interface ('iSCSI') technologies to provide block-level access to storage resources **308** within the storage system **306**, and other communications resources that that may be useful in facilitating data communications between components within the storage system **306**, as well as data communications between the storage system **306** and computing devices that are outside of the storage system **306**.

(96) The storage system **306** depicted in FIG. **3B** also includes processing resources **312** that may be useful in useful in executing computer program instructions and performing other computational tasks within the storage system **306**. The processing resources **312** may include one or more ASICs that are customized for some particular purpose as well as one or more CPUs. The processing resources **312** may also include one or more DSPs, one or more FPGAs, one or more systems on a chip ('SoCs'), or other form of processing resources **312**. The storage system **306** may utilize the storage resources **312** to perform a variety of tasks including, but not limited to, supporting the execution of software resources **314** that will be described in greater detail below.

(97) The storage system **306** depicted in FIG. **3B** also includes software resources **314** that, when executed by processing resources **312** within the storage system **306**, may perform a vast array of tasks. The software resources **314** may include, for example, one or more modules of computer program instructions that when executed by processing resources **312** within the storage system **306** are useful in carrying out various data protection techniques to preserve the integrity of data that is stored within the storage systems. Readers will appreciate that such data protection techniques may be carried out, for example, by system software executing on computer hardware within the storage system, by a cloud services provider, or in other ways. Such data protection

techniques can include, for example, data archiving techniques that cause data that is no longer actively used to be moved to a separate storage device or separate storage system for long-term retention, data backup techniques through which data stored in the storage system may be copied and stored in a distinct location to avoid data loss in the event of equipment failure or some other form of catastrophe with the storage system, data replication techniques through which data stored in the storage system is replicated to another storage system such that the data may be accessible via multiple storage systems, data snapshotting techniques through which the state of data within the storage system is captured at various points in time, data and database cloning techniques through which duplicate copies of data and databases may be created, and other data protection techniques.

(98) The software resources **314** may also include software that is useful in implementing software-defined storage ('SDS'). In such an example, the software resources **314** may include one or more modules of computer program instructions that, when executed, are useful in policy-based provisioning and management of data storage that is independent of the underlying hardware. Such software resources **314** may be useful in implementing storage virtualization to separate the storage hardware from the software that manages the storage hardware.

(99) The software resources **314** may also include software that is useful in facilitating and optimizing I/O operations that are directed to the storage resources **308** in the storage system **306**. For example, the software resources **314** may include software modules that perform carry out various data reduction techniques such as, for example, data compression, data deduplication, and others. The software resources **314** may include software modules that intelligently group together I/O operations to facilitate better usage of the underlying storage resource **308**, software modules that perform data migration operations to migrate from within a storage system, as well as software modules that perform other functions. Such software resources **314** may be embodied as one or more software containers or in many other ways.

(100) For further explanation, FIG. 3C sets forth an example of a cloud-based storage system **318** in accordance with some embodiments of the present disclosure. In the example depicted in FIG. 3C, the cloud-based storage system **318** is created entirely in a cloud computing environment **316** such as, for example, Amazon Web Services ('AWS'), Microsoft Azure, Google Cloud Platform, IBM Cloud, Oracle Cloud, and others. The cloud-based storage system **318** may be used to provide services similar to the services that may be provided by the storage systems described above. For example, the cloud-based storage system **318** may be used to provide block storage services to users of the cloud-based storage system **318**, the cloud-based storage system **318** may be used to provide storage services to users of the cloud-based storage system **318** through the use of solid-state storage, and so on.

(101) The cloud-based storage system **318** depicted in FIG. 3C includes two cloud computing instances **320**, **322** that each are used to support the execution of a storage controller application **324**, **326**. The cloud computing instances **320**, **322** may be embodied, for example, as instances of cloud computing resources (e.g., virtual machines) that may be provided by the cloud computing environment **316** to support the execution of software applications such as the storage controller application **324**, **326**. In one embodiment, the cloud computing instances **320**, **322** may be embodied as Amazon Elastic Compute Cloud ('EC2') instances. In such an example, an Amazon Machine Image ('AMI') that includes the storage controller application **324**, **326** may be booted to create and configure a virtual machine that may execute the storage controller application **324**, **326**.

(102) In the example method depicted in FIG. 3C, the storage controller application **324**, **326** may be embodied as a module of computer program instructions that, when executed, carries out various storage tasks. For example, the storage controller application **324**, **326** may be embodied as a module of computer program instructions that, when executed, carries out the same tasks as the controllers **110A**, **110B** in FIG. 1A described above such as writing data received from the users of the cloud-based storage system **318** to the cloud-based storage system **318**, erasing data from the

cloud-based storage system **318**, retrieving data from the cloud-based storage system **318** and providing such data to users of the cloud-based storage system **318**, monitoring and reporting of disk utilization and performance, performing redundancy operations, such as RAID or RAID-like data redundancy operations, compressing data, encrypting data, deduplicating data, and so forth. Readers will appreciate that because there are two cloud computing instances **320**, **322** that each include the storage controller application **324**, **326**, in some embodiments one cloud computing instance **320** may operate as the primary controller as described above while the other cloud computing instance **322** may operate as the secondary controller as described above. Readers will appreciate that the storage controller application **324**, **326** depicted in FIG. 3C may include identical source code that is executed within different cloud computing instances **320**, **322**.

(103) Consider an example in which the cloud computing environment **316** is embodied as AWS and the cloud computing instances are embodied as EC2 instances. In such an example, the cloud computing instance **320** that operates as the primary controller may be deployed on one of the instance types that has a relatively large amount of memory and processing power while the cloud computing instance **322** that operates as the secondary controller may be deployed on one of the instance types that has a relatively small amount of memory and processing power. In such an example, upon the occurrence of a failover event where the roles of primary and secondary are switched, a double failover may actually be carried out such that: 1) a first failover event where the cloud computing instance **322** that formerly operated as the secondary controller begins to operate as the primary controller, and 2) a third cloud computing instance (not shown) that is of an instance type that has a relatively large amount of memory and processing power is spun up with a copy of the storage controller application, where the third cloud computing instance begins operating as the primary controller while the cloud computing instance **322** that originally operated as the secondary controller begins operating as the secondary controller again. In such an example, the cloud computing instance **320** that formerly operated as the primary controller may be terminated. Readers will appreciate that in alternative embodiments, the cloud computing instance **320** that is operating as the secondary controller after the failover event may continue to operate as the secondary controller and the cloud computing instance **322** that operated as the primary controller after the occurrence of the failover event may be terminated once the primary role has been assumed by the third cloud computing instance (not shown).

(104) Readers will appreciate that while the embodiments described above relate to embodiments where one cloud computing instance **320** operates as the primary controller and the second cloud computing instance **322** operates as the secondary controller, other embodiments are within the scope of the present disclosure. For example, each cloud computing instance **320**, **322** may operate as a primary controller for some portion of the address space supported by the cloud-based storage system **318**, each cloud computing instance **320**, **322** may operate as a primary controller where the servicing of I/O operations directed to the cloud-based storage system **318** are divided in some other way, and so on. In fact, in other embodiments where costs savings may be prioritized over performance demands, only a single cloud computing instance may exist that contains the storage controller application.

(105) The cloud-based storage system **318** depicted in FIG. 3C includes cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338**. The cloud computing instances **340a**, **340b**, **340n** depicted in FIG. 3C may be embodied, for example, as instances of cloud computing resources that may be provided by the cloud computing environment **316** to support the execution of software applications. The cloud computing instances **340a**, **340b**, **340n** of FIG. 3C may differ from the cloud computing instances **320**, **322** described above as the cloud computing instances **340a**, **340b**, **340n** of FIG. 3C have local storage **330**, **334**, **338** resources whereas the cloud computing instances **320**, **322** that support the execution of the storage controller application **324**, **326** need not have local storage resources. The cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** may be embodied, for example, as EC2 M5 instances that include one

or more SSDs, as EC2 R5 instances that include one or more SSDs, as EC2 I3 instances that include one or more SSDs, and so on. In some embodiments, the local storage **330, 334, 338** must be embodied as solid-state storage (e.g., SSDs) rather than storage that makes use of hard disk drives.

(106) In the example depicted in FIG. 3C, each of the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338** can include a software daemon **328, 332, 336** that, when executed by a cloud computing instance **340a, 340b, 340n** can present itself to the storage controller applications **324, 326** as if the cloud computing instance **340a, 340b, 340n** were a physical storage device (e.g., one or more SSDs). In such an example, the software daemon **328, 332, 336** may include computer program instructions similar to those that would normally be contained on a storage device such that the storage controller applications **324, 326** can send and receive the same commands that a storage controller would send to storage devices. In such a way, the storage controller applications **324, 326** may include code that is identical to (or substantially identical to) the code that would be executed by the controllers in the storage systems described above. In these and similar embodiments, communications between the storage controller applications **324, 326** and the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338** may utilize iSCSI, NVMe over TCP, messaging, a custom protocol, or in some other mechanism.

(107) In the example depicted in FIG. 3C, each of the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338** may also be coupled to block-storage **342, 344, 346** that is offered by the cloud computing environment **316**. The block-storage **342, 344, 346** that is offered by the cloud computing environment **316** may be embodied, for example, as Amazon Elastic Block Store ('EBS') volumes. For example, a first EBS volume may be coupled to a first cloud computing instance **340a**, a second EBS volume may be coupled to a second cloud computing instance **340b**, and a third EBS volume may be coupled to a third cloud computing instance **340n**. In such an example, the block-storage **342, 344, 346** that is offered by the cloud computing environment **316** may be utilized in a manner that is similar to how the NVRAM devices described above are utilized, as the software daemon **328, 332, 336** (or some other module) that is executing within a particular cloud computing instance **340a, 340b, 340n** may, upon receiving a request to write data, initiate a write of the data to its attached EBS volume as well as a write of the data to its local storage **330, 334, 338** resources. In some alternative embodiments, data may only be written to the local storage **330, 334, 338** resources within a particular cloud computing instance **340a, 340b, 340n**. In an alternative embodiment, rather than using the block-storage **342, 344, 346** that is offered by the cloud computing environment **316** as NVRAM, actual RAM on each of the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338** may be used as NVRAM, thereby decreasing network utilization costs that would be associated with using an EBS volume as the NVRAM.

(108) In the example depicted in FIG. 3C, the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338** may be utilized, by cloud computing instances **320, 322** that support the execution of the storage controller application **324, 326** to service I/O operations that are directed to the cloud-based storage system **318**. Consider an example in which a first cloud computing instance **320** that is executing the storage controller application **324** is operating as the primary controller. In such an example, the first cloud computing instance **320** that is executing the storage controller application **324** may receive (directly or indirectly via the secondary controller) requests to write data to the cloud-based storage system **318** from users of the cloud-based storage system **318**. In such an example, the first cloud computing instance **320** that is executing the storage controller application **324** may perform various tasks such as, for example, deduplicating the data contained in the request, compressing the data contained in the request, determining where to write the data contained in the request, and so on, before ultimately sending a request to write a deduplicated, encrypted, or otherwise possibly updated version of the data to one or more of the cloud computing instances **340a, 340b, 340n** with local storage **330, 334, 338**. Either cloud

computing instance **320**, **322**, in some embodiments, may receive a request to read data from the cloud-based storage system **318** and may ultimately send a request to read data to one or more of the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338**.

(109) Readers will appreciate that when a request to write data is received by a particular cloud computing instance **340a**, **340b**, **340n** with local storage **330**, **334**, **338**, the software daemon **328**, **332**, **336** or some other module of computer program instructions that is executing on the particular cloud computing instance **340a**, **340b**, **340n** may be configured to not only write the data to its own local storage **330**, **334**, **338** resources and any appropriate block-storage **342**, **344**, **346** that are offered by the cloud computing environment **316**, but the software daemon **328**, **332**, **336** or some other module of computer program instructions that is executing on the particular cloud computing instance **340a**, **340b**, **340n** may also be configured to write the data to cloud-based object storage **348** that is attached to the particular cloud computing instance **340a**, **340b**, **340n**. The cloud-based object storage **348** that is attached to the particular cloud computing instance **340a**, **340b**, **340n** may be embodied, for example, as Amazon Simple Storage Service ('S3') storage that is accessible by the particular cloud computing instance **340a**, **340b**, **340n**. In other embodiments, the cloud computing instances **320**, **322** that each include the storage controller application **324**, **326** may initiate the storage of the data in the local storage **330**, **334**, **338** of the cloud computing instances **340a**, **340b**, **340n** and the cloud-based object storage **348**.

(110) Readers will appreciate that, as described above, the cloud-based storage system **318** may be used to provide block storage services to users of the cloud-based storage system **318**. While the local storage **330**, **334**, **338** resources and the block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n** may support block-level access, the cloud-based object storage **348** that is attached to the particular cloud computing instance **340a**, **340b**, **340n** supports only object-based access. In order to address this, the software daemon **328**, **332**, **336** or some other module of computer program instructions that is executing on the particular cloud computing instance **340a**, **340b**, **340n** may be configured to take blocks of data, package those blocks into objects, and write the objects to the cloud-based object storage **348** that is attached to the particular cloud computing instance **340a**, **340b**, **340n**.

(111) Consider an example in which data is written to the local storage **330**, **334**, **338** resources and the block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n** in 1 MB blocks. In such an example, assume that a user of the cloud-based storage system **318** issues a request to write data that, after being compressed and deduplicated by the storage controller application **324**, **326** results in the need to write 5 MB of data. In such an example, writing the data to the local storage **330**, **334**, **338** resources and the block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n** is relatively straightforward as 5 blocks that are 1 MB in size are written to the local storage **330**, **334**, **338** resources and the block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**. In such an example, the software daemon **328**, **332**, **336** or some other module of computer program instructions that is executing on the particular cloud computing instance **340a**, **340b**, **340n** may be configured to: 1) create a first object that includes the first 1 MB of data and write the first object to the cloud-based object storage **348**, 2) create a second object that includes the second 1 MB of data and write the second object to the cloud-based object storage **348**, 3) create a third object that includes the third 1 MB of data and write the third object to the cloud-based object storage **348**, and so on. As such, in some embodiments, each object that is written to the cloud-based object storage **348** may be identical (or nearly identical) in size. Readers will appreciate that in such an example, metadata that is associated with the data itself may be included in each object (e.g., the first 1 MB of the object is data and the remaining portion is metadata associated with the data).

(112) Readers will appreciate that the cloud-based object storage **348** may be incorporated into the cloud-based storage system **318** to increase the durability of the cloud-based storage system **318**.

Continuing with the example described above where the cloud computing instances **340a**, **340b**, **340n** are EC2 instances, readers will understand that EC2 instances are only guaranteed to have a monthly uptime of 99.9% and data stored in the local instance store only persists during the lifetime of the EC2 instance. As such, relying on the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** as the only source of persistent data storage in the cloud-based storage system **318** may result in a relatively unreliable storage system. Likewise, EBS volumes are designed for 99.999% availability. As such, even relying on EBS as the persistent data store in the cloud-based storage system **318** may result in a storage system that is not sufficiently durable. Amazon S3, however, is designed to provide 99.999999999% durability, meaning that a cloud-based storage system **318** that can incorporate S3 into its pool of storage is substantially more durable than various other options.

(113) Readers will appreciate that while a cloud-based storage system **318** that can incorporate S3 into its pool of storage is substantially more durable than various other options, utilizing S3 as the primary pool of storage may result in storage system that has relatively slow response times and relatively long I/O latencies. As such, the cloud-based storage system **318** depicted in FIG. 3C not only stores data in S3 but the cloud-based storage system **318** also stores data in local storage **330**, **334**, **338** resources and block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**, such that read operations can be serviced from local storage **330**, **334**, **338** resources and the block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**, thereby reducing read latency when users of the cloud-based storage system **318** attempt to read data from the cloud-based storage system **318**.

(114) In some embodiments, all data that is stored by the cloud-based storage system **318** may be stored in both: 1) the cloud-based object storage **348**, and 2) at least one of the local storage **330**, **334**, **338** resources or block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**. In such embodiments, the local storage **330**, **334**, **338** resources and block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n** may effectively operate as cache that generally includes all data that is also stored in S3, such that all reads of data may be serviced by the cloud computing instances **340a**, **340b**, **340n** without requiring the cloud computing instances **340a**, **340b**, **340n** to access the cloud-based object storage **348**. Readers will appreciate that in other embodiments, however, all data that is stored by the cloud-based storage system **318** may be stored in the cloud-based object storage **348**, but less than all data that is stored by the cloud-based storage system **318** may be stored in at least one of the local storage **330**, **334**, **338** resources or block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**. In such an example, various policies may be utilized to determine which subset of the data that is stored by the cloud-based storage system **318** should reside in both: 1) the cloud-based object storage **348**, and 2) at least one of the local storage **330**, **334**, **338** resources or block-storage **342**, **344**, **346** resources that are utilized by the cloud computing instances **340a**, **340b**, **340n**.

(115) As described above, when the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** are embodied as EC2 instances, the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** are only guaranteed to have a monthly uptime of 99.9% and data stored in the local instance store only persists during the lifetime of each cloud computing instance **340a**, **340b**, **340n** with local storage **330**, **334**, **338**. As such, one or more modules of computer program instructions that are executing within the cloud-based storage system **318** (e.g., a monitoring module that is executing on its own EC2 instance) may be designed to handle the failure of one or more of the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338**. In such an example, the monitoring module may handle the failure of one or more of the cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** by creating one or more new cloud computing instances with local storage, retrieving data that was stored on the failed cloud computing instances **340a**, **340b**, **340n** from the cloud-based object storage **348**, and

storing the data retrieved from the cloud-based object storage **348** in local storage on the newly created cloud computing instances. Readers will appreciate that many variants of this process may be implemented.

(116) Consider an example in which all cloud computing instances **340a**, **340b**, **340n** with local storage **330**, **334**, **338** failed. In such an example, the monitoring module may create new cloud computing instances with local storage, where high-bandwidth instances types are selected that allow for the maximum data transfer rates between the newly created high-bandwidth cloud computing instances with local storage and the cloud-based object storage **348**. Readers will appreciate that instances types are selected that allow for the maximum data transfer rates between the new cloud computing instances and the cloud-based object storage **348** such that the new high-bandwidth cloud computing instances can be rehydrated with data from the cloud-based object storage **348** as quickly as possible. Once the new high-bandwidth cloud computing instances are rehydrated with data from the cloud-based object storage **348**, less expensive lower-bandwidth cloud computing instances may be created, data may be migrated to the less expensive lower-bandwidth cloud computing instances, and the high-bandwidth cloud computing instances may be terminated.

(117) Readers will appreciate that in some embodiments, the number of new cloud computing instances that are created may substantially exceed the number of cloud computing instances that are needed to locally store all of the data stored by the cloud-based storage system **318**. The number of new cloud computing instances that are created may substantially exceed the number of cloud computing instances that are needed to locally store all of the data stored by the cloud-based storage system **318** in order to more rapidly pull data from the cloud-based object storage **348** and into the new cloud computing instances, as each new cloud computing instance can (in parallel) retrieve some portion of the data stored by the cloud-based storage system **318**. In such embodiments, once the data stored by the cloud-based storage system **318** has been pulled into the newly created cloud computing instances, the data may be consolidated within a subset of the newly created cloud computing instances and those newly created cloud computing instances that are excessive may be terminated.

(118) Consider an example in which 1000 cloud computing instances are needed in order to locally store all valid data that users of the cloud-based storage system **318** have written to the cloud-based storage system **318**. In such an example, assume that all 1,000 cloud computing instances fail. In such an example, the monitoring module may cause 100,000 cloud computing instances to be created, where each cloud computing instance is responsible for retrieving, from the cloud-based object storage **348**, distinct 1/100,000th chunks of the valid data that users of the cloud-based storage system **318** have written to the cloud-based storage system **318** and locally storing the distinct chunk of the dataset that it retrieved. In such an example, because each of the 100,000 cloud computing instances can retrieve data from the cloud-based object storage **348** in parallel, the caching layer may be restored 100 times faster as compared to an embodiment where the monitoring module only create 1000 replacement cloud computing instances. In such an example, over time the data that is stored locally in the 100,000 could be consolidated into 1,000 cloud computing instances and the remaining 99,000 cloud computing instances could be terminated.

(119) Readers will appreciate that various performance aspects of the cloud-based storage system **318** may be monitored (e.g., by a monitoring module that is executing in an EC2 instance) such that the cloud-based storage system **318** can be scaled-up or scaled-out as needed. Consider an example in which the monitoring module monitors the performance of the cloud-based storage system **318** via communications with one or more of the cloud computing instances **320**, **322** that each are used to support the execution of a storage controller application **324**, **326**, via monitoring communications between cloud computing instances **320**, **322**, **340a**, **340b**, **340n**, via monitoring communications between cloud computing instances **320**, **322**, **340a**, **340b**, **340n** and the cloud-based object storage **348**, or in some other way. In such an example, assume that the monitoring

module determines that the cloud computing instances **320**, **322** that are used to support the execution of a storage controller application **324**, **326** are undersized and not sufficiently servicing the I/O requests that are issued by users of the cloud-based storage system **318**. In such an example, the monitoring module may create a new, more powerful cloud computing instance (e.g., a cloud computing instance of a type that includes more processing power, more memory, etc.) that includes the storage controller application such that the new, more powerful cloud computing instance can begin operating as the primary controller. Likewise, if the monitoring module determines that the cloud computing instances **320**, **322** that are used to support the execution of a storage controller application **324**, **326** are oversized and that cost savings could be gained by switching to a smaller, less powerful cloud computing instance, the monitoring module may create a new, less powerful (and less expensive) cloud computing instance that includes the storage controller application such that the new, less powerful cloud computing instance can begin operating as the primary controller.

(120) Consider, as an additional example of dynamically sizing the cloud-based storage system **318**, an example in which the monitoring module determines that the utilization of the local storage that is collectively provided by the cloud computing instances **340a**, **340b**, **340n** has reached a predetermined utilization threshold (e.g., 95%). In such an example, the monitoring module may create additional cloud computing instances with local storage to expand the pool of local storage that is offered by the cloud computing instances. Alternatively, the monitoring module may create one or more new cloud computing instances that have larger amounts of local storage than the already existing cloud computing instances **340a**, **340b**, **340n**, such that data stored in an already existing cloud computing instance **340a**, **340b**, **340n** can be migrated to the one or more new cloud computing instances and the already existing cloud computing instance **340a**, **340b**, **340n** can be terminated, thereby expanding the pool of local storage that is offered by the cloud computing instances. Likewise, if the pool of local storage that is offered by the cloud computing instances is unnecessarily large, data can be consolidated and some cloud computing instances can be terminated.

(121) Readers will appreciate that the cloud-based storage system **318** may be sized up and down automatically by a monitoring module applying a predetermined set of rules that may be relatively simple or relatively complicated. In fact, the monitoring module may not only take into account the current state of the cloud-based storage system **318**, but the monitoring module may also apply predictive policies that are based on, for example, observed behavior (e.g., every night from 10 PM until 6 AM usage of the storage system is relatively light), predetermined fingerprints (e.g., every time a virtual desktop infrastructure adds 100 virtual desktops, the number of IOPS directed to the storage system increase by X), and so on. In such an example, the dynamic scaling of the cloud-based storage system **318** may be based on current performance metrics, predicted workloads, and many other factors, including combinations thereof.

(122) Readers will further appreciate that because the cloud-based storage system **318** may be dynamically scaled, the cloud-based storage system **318** may even operate in a way that is more dynamic. Consider the example of garbage collection. In a traditional storage system, the amount of storage is fixed. As such, at some point the storage system may be forced to perform garbage collection as the amount of available storage has become so constrained that the storage system is on the verge of running out of storage. In contrast, the cloud-based storage system **318** described here can always 'add' additional storage (e.g., by adding more cloud computing instances with local storage). Because the cloud-based storage system **318** described here can always 'add' additional storage, the cloud-based storage system **318** can make more intelligent decisions regarding when to perform garbage collection. For example, the cloud-based storage system **318** may implement a policy that garbage collection only be performed when the number of IOPS being serviced by the cloud-based storage system **318** falls below a certain level. In some embodiments, other system-level functions (e.g., deduplication, compression) may also be turned off and on in

response to system load, given that the size of the cloud-based storage system **318** is not constrained in the same way that traditional storage systems are constrained.

(123) Readers will appreciate that embodiments of the present disclosure resolve an issue with block-storage services offered by some cloud computing environments as some cloud computing environments only allow for one cloud computing instance to connect to a block-storage volume at a single time. For example, in Amazon AWS, only a single EC2 instance may be connected to an EBS volume. Through the use of EC2 instances with local storage, embodiments of the present disclosure can offer multi-connect capabilities where multiple EC2 instances can connect to another EC2 instance with local storage ('a drive instance'). In such embodiments, the drive instances may include software executing within the drive instance that allows the drive instance to support I/O directed to a particular volume from each connected EC2 instance. As such, some embodiments of the present disclosure may be embodied as multi-connect block storage services that may not include all of the components depicted in FIG. 3C.

(124) In some embodiments, especially in embodiments where the cloud-based object storage **348** resources are embodied as Amazon S3, the cloud-based storage system **318** may include one or more modules (e.g., a module of computer program instructions executing on an EC2 instance) that are configured to ensure that when the local storage of a particular cloud computing instance is rehydrated with data from S3, the appropriate data is actually in S3. This issue arises largely because S3 implements an eventual consistency model where, when overwriting an existing object, reads of the object will eventually (but not necessarily immediately) become consistent and will eventually (but not necessarily immediately) return the overwritten version of the object. To address this issue, in some embodiments of the present disclosure, objects in S3 are never overwritten. Instead, a traditional 'overwrite' would result in the creation of the new object (that includes the updated version of the data) and the eventual deletion of the old object (that includes the previous version of the data).

(125) In some embodiments of the present disclosure, as part of an attempt to never (or almost never) overwrite an object, when data is written to S3 the resultant object may be tagged with a sequence number. In some embodiments, these sequence numbers may be persisted elsewhere (e.g., in a database) such that at any point in time, the sequence number associated with the most up-to-date version of some piece of data can be known. In such a way, a determination can be made as to whether S3 has the most recent version of some piece of data by merely reading the sequence number associated with an object—and without actually reading the data from S3. The ability to make this determination may be particularly important when a cloud computing instance with local storage crashes, as it would be undesirable to rehydrate the local storage of a replacement cloud computing instance with out-of-date data. In fact, because the cloud-based storage system **318** does not need to access the data to verify its validity, the data can stay encrypted and access charges can be avoided.

(126) The storage systems described above may carry out intelligent data backup techniques through which data stored in the storage system may be copied and stored in a distinct location to avoid data loss in the event of equipment failure or some other form of catastrophe. For example, the storage systems described above may be configured to examine each backup to avoid restoring the storage system to an undesirable state. Consider an example in which malware infects the storage system. In such an example, the storage system may include software resources **314** that can scan each backup to identify backups that were captured before the malware infected the storage system and those backups that were captured after the malware infected the storage system. In such an example, the storage system may restore itself from a backup that does not include the malware—or at least not restore the portions of a backup that contained the malware. In such an example, the storage system may include software resources **314** that can scan each backup to identify the presences of malware (or a virus, or some other undesirable), for example, by identifying write operations that were serviced by the storage system and originated from a network

subnet that is suspected to have delivered the malware, by identifying write operations that were serviced by the storage system and originated from a user that is suspected to have delivered the malware, by identifying write operations that were serviced by the storage system and examining the content of the write operation against fingerprints of the malware, and in many other ways.

(127) Readers will further appreciate that the backups (often in the form of one or more snapshots) may also be utilized to perform rapid recovery of the storage system. Consider an example in which the storage system is infected with ransomware that locks users out of the storage system. In such an example, software resources **314** within the storage system may be configured to detect the presence of ransomware and may be further configured to restore the storage system to a point-in-time, using the retained backups, prior to the point-in-time at which the ransomware infected the storage system. In such an example, the presence of ransomware may be explicitly detected through the use of software tools utilized by the system, through the use of a key (e.g., a USB drive) that is inserted into the storage system, or in a similar way. Likewise, the presence of ransomware may be inferred in response to system activity meeting a predetermined fingerprint such as, for example, no reads or writes coming into the system for a predetermined period of time.

(128) Readers will appreciate that the various components described above may be grouped into one or more optimized computing packages as converged infrastructures. Such converged infrastructures may include pools of computers, storage and networking resources that can be shared by multiple applications and managed in a collective manner using policy-driven processes. Such converged infrastructures may be implemented with a converged infrastructure reference architecture, with standalone appliances, with a software driven hyper-converged approach (e.g., hyper-converged infrastructures), or in other ways.

(129) Readers will appreciate that the storage systems described above may be useful for supporting various types of software applications. For example, the storage system **306** may be useful in supporting artificial intelligence ('AI') applications, database applications, DevOps projects, electronic design automation tools, event-driven software applications, high performance computing applications, simulation applications, high-speed data capture and analysis applications, machine learning applications, media production applications, media serving applications, picture archiving and communication systems ('PACS') applications, software development applications, virtual reality applications, augmented reality applications, and many other types of applications by providing storage resources to such applications.

(130) The storage systems described above may operate to support a wide variety of applications. In view of the fact that the storage systems include compute resources, storage resources, and a wide variety of other resources, the storage systems may be well suited to support applications that are resource intensive such as, for example, AI applications. AI applications may be deployed in a variety of fields, including: predictive maintenance in manufacturing and related fields, healthcare applications such as patient data & risk analytics, retail and marketing deployments (e.g., search advertising, social media advertising), supply chains solutions, fintech solutions such as business analytics & reporting tools, operational deployments such as real-time analytics tools, application performance management tools, IT infrastructure management tools, and many others.

(131) Such AI applications may enable devices to perceive their environment and take actions that maximize their chance of success at some goal. Examples of such AI applications can include IBM Watson, Microsoft Oxford, Google DeepMind, Baidu Minwa, and others. The storage systems described above may also be well suited to support other types of applications that are resource intensive such as, for example, machine learning applications. Machine learning applications may perform various types of data analysis to automate analytical model building. Using algorithms that iteratively learn from data, machine learning applications can enable computers to learn without being explicitly programmed. One particular area of machine learning is referred to as reinforcement learning, which involves taking suitable actions to maximize reward in a particular situation. Reinforcement learning may be employed to find the best possible behavior or path that a

particular software application or machine should take in a specific situation. Reinforcement learning differs from other areas of machine learning (e.g., supervised learning, unsupervised learning) in that correct input/output pairs need not be presented for reinforcement learning and sub-optimal actions need not be explicitly corrected.

(132) In addition to the resources already described, the storage systems described above may also include graphics processing units ('GPUs'), occasionally referred to as visual processing unit ('VPUs'). Such GPUs may be embodied as specialized electronic circuits that rapidly manipulate and alter memory to accelerate the creation of images in a frame buffer intended for output to a display device. Such GPUs may be included within any of the computing devices that are part of the storage systems described above, including as one of many individually scalable components of a storage system, where other examples of individually scalable components of such storage system can include storage components, memory components, compute components (e.g., CPUs, FPGAs, ASICs), networking components, software components, and others. In addition to GPUs, the storage systems described above may also include neural network processors ('NNPs') for use in various aspects of neural network processing. Such NNPs may be used in place of (or in addition to) GPUs and may also be independently scalable.

(133) As described above, the storage systems described herein may be configured to support artificial intelligence applications, machine learning applications, big data analytics applications, and many other types of applications. The rapid growth in these sort of applications is being driven by three technologies: deep learning (DL), GPU processors, and Big Data. Deep learning is a computing model that makes use of massively parallel neural networks inspired by the human brain. Instead of experts handcrafting software, a deep learning model writes its own software by learning from lots of examples. Such GPUs may include thousands of cores that are well-suited to run algorithms that loosely represent the parallel nature of the human brain.

(134) Advances in deep neural networks have ignited a new wave of algorithms and tools for data scientists to tap into their data with artificial intelligence (AI). With improved algorithms, larger data sets, and various frameworks (including open-source software libraries for machine learning across a range of tasks), data scientists are tackling new use cases like autonomous driving vehicles, natural language processing and understanding, computer vision, machine reasoning, strong AI, and many others. Applications of such techniques may include: machine and vehicular object detection, identification and avoidance; visual recognition, classification and tagging; algorithmic financial trading strategy performance management; simultaneous localization and mapping; predictive maintenance of high-value machinery; prevention against cyber security threats, expertise automation; image recognition and classification; question answering; robotics; text analytics (extraction, classification) and text generation and translation; and many others. Applications of AI techniques has materialized in a wide array of products include, for example, Amazon Echo's speech recognition technology that allows users to talk to their machines, Google Translate™ which allows for machine-based language translation, Spotify's Discover Weekly that provides recommendations on new songs and artists that a user may like based on the user's usage and traffic analysis, Quill's text generation offering that takes structured data and turns it into narrative stories, Chatbots that provide real-time, contextually specific answers to questions in a dialog format, and many others.

(135) Data is the heart of modern AI and deep learning algorithms. Before training can begin, one problem that must be addressed revolves around collecting the labeled data that is crucial for training an accurate AI model. A full scale AI deployment may be required to continuously collect, clean, transform, label, and store large amounts of data. Adding additional high quality data points directly translates to more accurate models and better insights. Data samples may undergo a series of processing steps including, but not limited to: 1) ingesting the data from an external source into the training system and storing the data in raw form, 2) cleaning and transforming the data in a format convenient for training, including linking data samples to the appropriate label, 3) exploring

parameters and models, quickly testing with a smaller dataset, and iterating to converge on the most promising models to push into the production cluster, 4) executing training phases to select random batches of input data, including both new and older samples, and feeding those into production GPU servers for computation to update model parameters, and 5) evaluating including using a holdback portion of the data not used in training in order to evaluate model accuracy on the holdout data. This lifecycle may apply for any type of parallelized machine learning, not just neural networks or deep learning. For example, standard machine learning frameworks may rely on CPUs instead of GPUs but the data ingest and training workflows may be the same. Readers will appreciate that a single shared storage data hub creates a coordination point throughout the lifecycle without the need for extra data copies among the ingest, preprocessing, and training stages. Rarely is the ingested data used for only one purpose, and shared storage gives the flexibility to train multiple different models or apply traditional analytics to the data.

(136) Readers will appreciate that each stage in the AI data pipeline may have varying requirements from the data hub (e.g., the storage system or collection of storage systems). Scale-out storage systems must deliver uncompromising performance for all manner of access types and patterns—from small, metadata-heavy to large files, from random to sequential access patterns, and from low to high concurrency. The storage systems described above may serve as an ideal AI data hub as the systems may service unstructured workloads. In the first stage, data is ideally ingested and stored on to the same data hub that following stages will use, in order to avoid excess data copying. The next two steps can be done on a standard compute server that optionally includes a GPU, and then in the fourth and last stage, full training production jobs are run on powerful GPU-accelerated servers. Often, there is a production pipeline alongside an experimental pipeline operating on the same dataset. Further, the GPU-accelerated servers can be used independently for different models or joined together to train on one larger model, even spanning multiple systems for distributed training. If the shared storage tier is slow, then data must be copied to local storage for each phase, resulting in wasted time staging data onto different servers. The ideal data hub for the AI training pipeline delivers performance similar to data stored locally on the server node while also having the simplicity and performance to enable all pipeline stages to operate concurrently.

(137) Although the preceding paragraphs discuss deep learning applications, readers will appreciate that the storage systems described herein may also be part of a distributed deep learning (‘DDL’) platform to support the execution of DDL algorithms. The storage systems described above may also be paired with other technologies such as TensorFlow, an open-source software library for dataflow programming across a range of tasks that may be used for machine learning applications such as neural networks, to facilitate the development of such machine learning models, applications, and so on.

(138) The storage systems described above may also be used in a neuromorphic computing environment. Neuromorphic computing is a form of computing that mimics brain cells. To support neuromorphic computing, an architecture of interconnected “neurons” replace traditional computing models with low-powered signals that go directly between neurons for more efficient computation. Neuromorphic computing may make use of very-large-scale integration (VLSI) systems containing electronic analog circuits to mimic neuro-biological architectures present in the nervous system, as well as analog, digital, mixed-mode analog/digital VLSI, and software systems that implement models of neural systems for perception, motor control, or multisensory integration.

(139) Readers will appreciate that the storage systems described above may be configured to support the storage or use of (among other types of data) blockchains. In addition to supporting the storage and use of blockchain technologies, the storage systems described above may also support the storage and use of derivative items such as, for example, open source blockchains and related tools that are part of the IBM™ Hyperledger project, permissioned blockchains in which a certain number of trusted parties are allowed to access the block chain, blockchain products that enable developers to build their own distributed ledger projects, and others. Blockchains and the storage

systems described herein may be leveraged to support on-chain storage of data as well as off-chain storage of data.

(140) Off-chain storage of data can be implemented in a variety of ways and can occur when the data itself is not stored within the blockchain. For example, in one embodiment, a hash function may be utilized and the data itself may be fed into the hash function to generate a hash value. In such an example, the hashes of large pieces of data may be embedded within transactions, instead of the data itself. Readers will appreciate that, in other embodiments, alternatives to blockchains may be used to facilitate the decentralized storage of information. For example, one alternative to a blockchain that may be used is a blockweave. While conventional blockchains store every transaction to achieve validation, a blockweave permits secure decentralization without the usage of the entire chain, thereby enabling low cost on-chain storage of data. Such blockweaves may utilize a consensus mechanism that is based on proof of access (PoA) and proof of work (PoW).

(141) The storage systems described above may, either alone or in combination with other computing devices, be used to support in-memory computing applications. In-memory computing involves the storage of information in RAM that is distributed across a cluster of computers. Readers will appreciate that the storage systems described above, especially those that are configurable with customizable amounts of processing resources, storage resources, and memory resources (e.g., those systems in which blades that contain configurable amounts of each type of resource), may be configured in a way so as to provide an infrastructure that can support in-memory computing. Likewise, the storage systems described above may include component parts (e.g., NVDIMMs, 3D crosspoint storage that provide fast random access memory that is persistent) that can actually provide for an improved in-memory computing environment as compared to in-memory computing environments that rely on RAM distributed across dedicated servers.

(142) In some embodiments, the storage systems described above may be configured to operate as a hybrid in-memory computing environment that includes a universal interface to all storage media (e.g., RAM, flash storage, 3D crosspoint storage). In such embodiments, users may have no knowledge regarding the details of where their data is stored but they can still use the same full, unified API to address data. In such embodiments, the storage system may (in the background) move data to the fastest layer available—including intelligently placing the data in dependence upon various characteristics of the data or in dependence upon some other heuristic. In such an example, the storage systems may even make use of existing products such as Apache Ignite and GridGain to move data between the various storage layers, or the storage systems may make use of custom software to move data between the various storage layers. The storage systems described herein may implement various optimizations to improve the performance of in-memory computing such as, for example, having computations occur as close to the data as possible.

(143) Readers will further appreciate that in some embodiments, the storage systems described above may be paired with other resources to support the applications described above. For example, one infrastructure could include primary compute in the form of servers and workstations which specialize in using General-purpose computing on graphics processing units (‘GPGPU’) to accelerate deep learning applications that are interconnected into a computation engine to train parameters for deep neural networks. Each system may have Ethernet external connectivity, InfiniBand external connectivity, some other form of external connectivity, or some combination thereof. In such an example, the GPUs can be grouped for a single large training or used independently to train multiple models. The infrastructure could also include a storage system such as those described above to provide, for example, a scale-out all-flash file or object store through which data can be accessed via high-performance protocols such as NFS, S3, and so on. The infrastructure can also include, for example, redundant top-of-rack Ethernet switches connected to storage and compute via ports in MLAG port channels for redundancy. The infrastructure could also include additional compute in the form of whitebox servers, optionally with GPUs, for data ingestion, pre-processing, and model debugging. Readers will appreciate that additional

infrastructures are also be possible.

(144) Readers will appreciate that the storage systems described above, either alone or in coordination with other computing machinery may be configured to support other AI related tools. For example, the storage systems may make use of tools like ONNX or other open neural network exchange formats that make it easier to transfer models written in different AI frameworks. Likewise, the storage systems may be configured to support tools like Amazon's Glue that allow developers to prototype, build, and train deep learning models. In fact, the storage systems described above may be part of a larger platform, such as IBM™ Cloud Private for Data, that includes integrated data science, data engineering and application building services.

(145) Readers will further appreciate that the storage systems described above may also be deployed as an edge solution. Such an edge solution may be in place to optimize cloud computing systems by performing data processing at the edge of the network, near the source of the data. Edge computing can push applications, data and computing power (i.e., services) away from centralized points to the logical extremes of a network. Through the use of edge solutions such as the storage systems described above, computational tasks may be performed using the compute resources provided by such storage systems, data may be storage using the storage resources of the storage system, and cloud-based services may be accessed through the use of various resources of the storage system (including networking resources). By performing computational tasks on the edge solution, storing data on the edge solution, and generally making use of the edge solution, the consumption of expensive cloud-based resources may be avoided and, in fact, performance improvements may be experienced relative to a heavier reliance on cloud-based resources.

(146) While many tasks may benefit from the utilization of an edge solution, some particular uses may be especially suited for deployment in such an environment. For example, devices like drones, autonomous cars, robots, and others may require extremely rapid processing—so fast, in fact, that sending data up to a cloud environment and back to receive data processing support may simply be too slow. As an additional example, some IoT devices such as connected video cameras may not be well-suited for the utilization of cloud-based resources as it may be impractical (not only from a privacy perspective, security perspective, or a financial perspective) to send the data to the cloud simply because of the pure volume of data that is involved. As such, many tasks that really on data processing, storage, or communications may be better suited by platforms that include edge solutions such as the storage systems described above.

(147) The storage systems described above may alone, or in combination with other computing resources, serves as a network edge platform that combines compute resources, storage resources, networking resources, cloud technologies and network virtualization technologies, and so on. As part of the network, the edge may take on characteristics similar to other network facilities, from the customer premise and backhaul aggregation facilities to Points of Presence (PoPs) and regional data centers. Readers will appreciate that network workloads, such as Virtual Network Functions (VNFs) and others, will reside on the network edge platform. Enabled by a combination of containers and virtual machines, the network edge platform may rely on controllers and schedulers that are no longer geographically co-located with the data processing resources. The functions, as microservices, may split into control planes, user and data planes, or even state machines, allowing for independent optimization and scaling techniques to be applied. Such user and data planes may be enabled through increased accelerators, both those residing in server platforms, such as FPGAs and Smart NICs, and through SDN-enabled merchant silicon and programmable ASICs.

(148) The storage systems described above may also be optimized for use in big data analytics. Big data analytics may be generally described as the process of examining large and varied data sets to uncover hidden patterns, unknown correlations, market trends, customer preferences and other useful information that can help organizations make more-informed business decisions. As part of that process, semi-structured and unstructured data such as, for example, internet clickstream data, web server logs, social media content, text from customer emails and survey responses, mobile-

phone call-detail records, IoT sensor data, and other data may be converted to a structured form.

(149) The storage systems described above may also support (including implementing as a system interface) applications that perform tasks in response to human speech. For example, the storage systems may support the execution intelligent personal assistant applications such as, for example, Amazon's Alexa, Apple Siri, Google Voice, Samsung Bixby, Microsoft Cortana, and others. While the examples described in the previous sentence make use of voice as input, the storage systems described above may also support chatbots, talkbots, chatterbots, or artificial conversational entities or other applications that are configured to conduct a conversation via auditory or textual methods. Likewise, the storage system may actually execute such an application to enable a user such as a system administrator to interact with the storage system via speech. Such applications are generally capable of voice interaction, music playback, making to-do lists, setting alarms, streaming podcasts, playing audiobooks, and providing weather, traffic, and other real time information, such as news, although in embodiments in accordance with the present disclosure, such applications may be utilized as interfaces to various system management operations.

(150) The storage systems described above may also implement AI platforms for delivering on the vision of self-driving storage. Such AI platforms may be configured to deliver global predictive intelligence by collecting and analyzing large amounts of storage system telemetry data points to enable effortless management, analytics and support. In fact, such storage systems may be capable of predicting both capacity and performance, as well as generating intelligent advice on workload deployment, interaction and optimization. Such AI platforms may be configured to scan all incoming storage system telemetry data against a library of issue fingerprints to predict and resolve incidents in real-time, before they impact customer environments, and captures hundreds of variables related to performance that are used to forecast performance load.

(151) The storage systems described above may support the serialized or simultaneous execution of artificial intelligence applications, machine learning applications, data analytics applications, data transformations, and other tasks that collectively may form an AI ladder. Such an AI ladder may effectively be formed by combining such elements to form a complete data science pipeline, where exist dependencies between elements of the AI ladder. For example, AI may require that some form of machine learning has taken place, machine learning may require that some form of analytics has taken place, analytics may require that some form of data and information architecting has taken place, and so on. As such, each element may be viewed as a rung in an AI ladder that collectively can form a complete and sophisticated AI solution.

(152) The storage systems described above may also, either alone or in combination with other computing environments, be used to deliver an AI everywhere experience where AI permeates wide and expansive aspects of business and life. For example, AI may play an important role in the delivery of deep learning solutions, deep reinforcement learning solutions, artificial general intelligence solutions, autonomous vehicles, cognitive computing solutions, commercial UAVs or drones, conversational user interfaces, enterprise taxonomies, ontology management solutions, machine learning solutions, smart dust, smart robots, smart workplaces, and many others.

(153) The storage systems described above may also, either alone or in combination with other computing environments, be used to deliver a wide range of transparently immersive experiences (including those that use digital twins of various “things” such as people, places, processes, systems, and so on) where technology can introduce transparency between people, businesses, and things. Such transparently immersive experiences may be delivered as augmented reality technologies, connected homes, virtual reality technologies, brain-computer interfaces, human augmentation technologies, nanotube electronics, volumetric displays, 4D printing technologies, or others.

(154) The storage systems described above may also, either alone or in combination with other computing environments, be used to support a wide variety of digital platforms. Such digital platforms can include, for example, 5G wireless systems and platforms, digital twin platforms,

edge computing platforms, IoT platforms, quantum computing platforms, serverless PaaS, software-defined security, neuromorphic computing platforms, and so on.

(155) The storage systems described above may also be part of a multi-cloud environment in which multiple cloud computing and storage services are deployed in a single heterogeneous architecture. In order to facilitate the operation of such a multi-cloud environment, DevOps tools may be deployed to enable orchestration across clouds. Likewise, continuous development and continuous integration tools may be deployed to standardize processes around continuous integration and delivery, new feature rollout and provisioning cloud workloads. By standardizing these processes, a multi-cloud strategy may be implemented that enables the utilization of the best provider for each workload.

(156) The storage systems described above may be used as a part of a platform to enable the use of crypto-anchors that may be used to authenticate a product's origins and contents to ensure that it matches a blockchain record associated with the product. Similarly, as part of a suite of tools to secure data stored on the storage system, the storage systems described above may implement various encryption technologies and schemes, including lattice cryptography. Lattice cryptography can involve constructions of cryptographic primitives that involve lattices, either in the construction itself or in the security proof. Unlike public-key schemes such as the RSA, Diffie-Hellman or Elliptic-Curve cryptosystems, which are easily attacked by a quantum computer, some lattice-based constructions appear to be resistant to attack by both classical and quantum computers.

(157) A quantum computer is a device that performs quantum computing. Quantum computing is computing using quantum-mechanical phenomena, such as superposition and entanglement. Quantum computers differ from traditional computers that are based on transistors, as such traditional computers require that data be encoded into binary digits (bits), each of which is always in one of two definite states (0 or 1). In contrast to traditional computers, quantum computers use quantum bits, which can be in superpositions of states. A quantum computer maintains a sequence of qubits, where a single qubit can represent a one, a zero, or any quantum superposition of those two qubit states. A pair of qubits can be in any quantum superposition of 4 states, and three qubits in any superposition of 8 states. A quantum computer with n qubits can generally be in an arbitrary superposition of up to 2^n different states simultaneously, whereas a traditional computer can only be in one of these states at any one time. A quantum Turing machine is a theoretical model of such a computer.

(158) The storage systems described above may also be paired with FPGA-accelerated servers as part of a larger AI or ML infrastructure. Such FPGA-accelerated servers may reside near (e.g., in the same data center) the storage systems described above or even incorporated into an appliance that includes one or more storage systems, one or more FPGA-accelerated servers, networking infrastructure that supports communications between the one or more storage systems and the one or more FPGA-accelerated servers, as well as other hardware and software components.

Alternatively, FPGA-accelerated servers may reside within a cloud computing environment that may be used to perform compute-related tasks for AI and ML jobs. Any of the embodiments described above may be used to collectively serve as a FPGA-based AI or ML platform. Readers will appreciate that, in some embodiments of the FPGA-based AI or ML platform, the FPGAs that are contained within the FPGA-accelerated servers may be reconfigured for different types of ML models (e.g., LSTMs, CNNs, GRUs). The ability to reconfigure the FPGAs that are contained within the FPGA-accelerated servers may enable the acceleration of a ML or AI application based on the most optimal numerical precision and memory model being used. Readers will appreciate that by treating the collection of FPGA-accelerated servers as a pool of FPGAs, any CPU in the data center may utilize the pool of FPGAs as a shared hardware microservice, rather than limiting a server to dedicated accelerators plugged into it.

(159) The FPGA-accelerated servers and the GPU-accelerated servers described above may

implement a model of computing where, rather than keeping a small amount of data in a CPU and running a long stream of instructions over it as occurred in more traditional computing models, the machine learning model and parameters are pinned into the high-bandwidth on-chip memory with lots of data streaming through the high-bandwidth on-chip memory. FPGAs may even be more efficient than GPUs for this computing model, as the FPGAs can be programmed with only the instructions needed to run this kind of computing model.

(160) The storage systems described above may be configured to provide parallel storage, for example, through the use of a parallel file system such as BeeGFS. Such parallel file systems may include a distributed metadata architecture. For example, the parallel file system may include a plurality of metadata servers across which metadata is distributed, as well as components that include services for clients and storage servers.

(161) The systems described above can support the execution of a wide array of software applications. Such software applications can be deployed in a variety of ways, including container-based deployment models. Containerized applications may be managed using a variety of tools. For example, containerized applications may be managed using Docker Swarm, Kubernetes, and others. Containerized applications may be used to facilitate a serverless, cloud native computing deployment and management model for software applications. In support of a serverless, cloud native computing deployment and management model for software applications, containers may be used as part of an event handling mechanisms (e.g., AWS Lambdas) such that various events cause a containerized application to be spun up to operate as an event handler.

(162) The systems described above may be deployed in a variety of ways, including being deployed in ways that support fifth generation ('5G') networks. 5G networks may support substantially faster data communications than previous generations of mobile communications networks and, as a consequence may lead to the disaggregation of data and computing resources as modern massive data centers may become less prominent and may be replaced, for example, by more-local, micro data centers that are close to the mobile-network towers. The systems described above may be included in such local, micro data centers and may be part of or paired to multi-access edge computing ('MEC') systems. Such MEC systems may enable cloud computing capabilities and an IT service environment at the edge of the cellular network. By running applications and performing related processing tasks closer to the cellular customer, network congestion may be reduced and applications may perform better.

(163) In some examples, a non-transitory computer-readable medium storing computer-readable instructions may be provided in accordance with the principles described herein. The instructions, when executed by a processor of a computing device, may direct the processor and/or computing device to perform one or more operations, including one or more of the operations described herein. Such instructions may be stored and/or transmitted using any of a variety of known computer-readable media.

(164) A non-transitory computer-readable medium as referred to herein may include any non-transitory storage medium that participates in providing data (e.g., instructions) that may be read and/or executed by a computing device (e.g., by a processor of a computing device). For example, a non-transitory computer-readable medium may include, but is not limited to, any combination of non-volatile storage media and/or volatile storage media. Exemplary non-volatile storage media include, but are not limited to, read-only memory, flash memory, a solid-state drive, a magnetic storage device (e.g. a hard disk, a floppy disk, magnetic tape, etc.), ferroelectric random-access memory ("RAM"), and an optical disc (e.g., a compact disc, a digital video disc, a Blu-ray disc, etc.). Exemplary volatile storage media include, but are not limited to, RAM (e.g., dynamic RAM).

(165) FIG. 3D illustrates an exemplary computing device 350 that may be specifically configured to perform one or more of the processes described herein. As shown in FIG. 3D, computing device 350 may include a communication interface 352, a processor 354, a storage device 356, and an input/output ("I/O") module 358 communicatively connected one to another via a communication

infrastructure **360**. While an exemplary computing device **350** is shown in FIG. 3D, the components illustrated in FIG. 3D are not intended to be limiting. Additional or alternative components may be used in other embodiments. Components of computing device **350** shown in FIG. 3D will now be described in additional detail.

(166) Communication interface **352** may be configured to communicate with one or more computing devices. Examples of communication interface **352** include, without limitation, a wired network interface (such as a network interface card), a wireless network interface (such as a wireless network interface card), a modem, an audio/video connection, and any other suitable interface.

(167) Processor **354** generally represents any type or form of processing unit capable of processing data and/or interpreting, executing, and/or directing execution of one or more of the instructions, processes, and/or operations described herein. Processor **354** may perform operations by executing computer-executable instructions **362** (e.g., an application, software, code, and/or other executable data instance) stored in storage device **356**.

(168) Storage device **356** may include one or more data storage media, devices, or configurations and may employ any type, form, and combination of data storage media and/or device. For example, storage device **356** may include, but is not limited to, any combination of the non-volatile media and/or volatile media described herein. Electronic data, including data described herein, may be temporarily and/or permanently stored in storage device **356**. For example, data representative of computer-executable instructions **362** configured to direct processor **354** to perform any of the operations described herein may be stored within storage device **356**. In some examples, data may be arranged in one or more databases residing within storage device **356**.

(169) I/O module **358** may include one or more I/O modules configured to receive user input and provide user output. I/O module **358** may include any hardware, firmware, software, or combination thereof supportive of input and output capabilities. For example, I/O module **358** may include hardware and/or software for capturing user input, including, but not limited to, a keyboard or keypad, a touchscreen component (e.g., touchscreen display), a receiver (e.g., an RF or infrared receiver), motion sensors, and/or one or more input buttons.

(170) I/O module **358** may include one or more devices for presenting output to a user, including, but not limited to, a graphics engine, a display (e.g., a display screen), one or more output drivers (e.g., display drivers), one or more audio speakers, and one or more audio drivers. In certain embodiments, I/O module **358** is configured to provide graphical data to a display for presentation to a user. The graphical data may be representative of one or more graphical user interfaces and/or any other graphical content as may serve a particular implementation. In some examples, any of the systems, computing devices, and/or other components described herein may be implemented by computing device **350**.

(171) Advantages and features of the present disclosure can be further described by the following statements. 1. A method comprising: identifying, by a data protection system, one or more attributes of a storage element within a storage system; determining, by the data protection system based on the one or more attributes of the storage element, a data protection policy for a dataset associated with the storage system; and applying, by the data protection system, the data protection policy to the dataset. 2. The method of any of the preceding statements, wherein the one or more attributes comprise one or more of a time of creation for the storage element, a purpose for which the storage element is created, a specified amount of time that the storage element is to be in existence before being eradicated, a type of data that the storage element is to store, an identity of a host associated with the storage element, an identity of a requestor that provides a request to create the storage element, an attribute of operations requested to be performed with respect to the storage element, or a tag associated with the storage element. 3. The method of any of the preceding statements, wherein the identifying the one or more attributes comprises accessing metadata associated with the storage element. 4. The method of any of the preceding statements, wherein the

accessing the metadata is performed in response to detecting a creation event in which the storage element is created. 5. The method of any of the preceding statements, wherein the determining the data protection policy comprises generating the data protection policy. 6. The method of any of the preceding statements, wherein the determining the data protection policy comprises updating the data protection policy. 7. The method of any of the preceding statements, wherein: the updating the data protection policy comprises causing the data protection policy to have a first ruleset to having a second ruleset that is less protective than the first ruleset; and the applying the data protection policy to the dataset comprises waiting, based on the second ruleset being less protective than the first ruleset, to apply the data protection policy to the dataset for a predetermined time delay. 8. The method of any of the preceding statements, wherein: the updating the data protection policy comprises causing the data protection policy to have a first ruleset to having a second ruleset that is less protective than the first ruleset; and the applying the data protection policy to the dataset comprises retaining, based on the second ruleset being less protective than the first ruleset, one or more rules from the first ruleset for inclusion in the second ruleset. 9. The method of any of the preceding statements, wherein: the identifying the one or more attributes comprises detecting a change in an attribute of the storage element; and the determining the data protection policy comprises updating, based on the change in the attribute of the storage element, the data protection policy. 10. The method of any of the preceding statements, wherein the determining the data protection policy is performed automatically without user input provided by a user associated with the storage system. 11. The method of any of the preceding statements, wherein the data protection policy specifies one or more data protection rules for one or more of the storage element, a recovery dataset generated by the storage system for the storage element, or the storage system. 12. The method of any of the preceding statements, further comprising: detecting, by the data protection system, a request to perform an operation with respect to the dataset; preventing, by the data protection system based on the data protection policy, the operation from being performed. 13. The method of any of the preceding statements, further comprising: detecting, by the data protection system, a request to perform an operation with respect to the dataset; allowing, by the data protection system based on the data protection policy, the operation to be performed. 14. The method of any of the preceding statements, further comprising: scanning the dataset subsequent to the applying the data protection policy to the dataset; determining, based on the scanning, that the data protection policy is potentially not appropriate for the dataset; and performing, based on the determining that the data protection policy is potentially not appropriate for the dataset, a predetermined action with respect to the data protection policy. 15. The method of any of the preceding statements, where the performing the predetermined action comprises providing an alert. 16. The method of any of the preceding statements, wherein the data protection system is implemented by a controller within the storage system. 17. The method of any of the preceding statements, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network. 18. The method of any of the preceding statements, wherein one or more of the identifying, the determining, or the applying is performed using a machine learning model. 19. A system comprising: a memory storing instructions; and one or more processors communicatively coupled to the memory and configured to execute the instructions to perform a process comprising: identifying, by a data protection system, one or more attributes of a storage element within a storage system; determining, by the data protection system based on the one or more attributes of the storage element, a data protection policy for a dataset associated with the storage system; and applying, by the data protection system, the data protection policy to the dataset. 20. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to perform a processing comprising: identifying, by a data protection system, one or more attributes of a storage element within a storage system; determining, by the data protection system based on the one or more attributes of the storage element, a data protection policy for a dataset associated with the storage

system; and applying, by the data protection system, the data protection policy to the dataset.

(172) Additional advantages and features of the present disclosure can be further described by the following statements. 1. A method comprising: determining, by a data protection system, that a total amount of read traffic and write traffic processed by a storage system during a time period exceeds a threshold, the read traffic representing data read from the storage system during the time period and the write traffic representing data written to the storage system during the time period; determining, by the data protection system, that the write traffic is less compressible than the read traffic; and determining, by the data protection system based on the total amount of read traffic and write traffic exceeding the threshold and on the write traffic being less compressible than the read traffic, that the storage system is possibly being targeted by a security threat. 2. The method of any of the preceding statements, further comprising: identifying, by the data protection system, an attribute associated with one or more of the data read from the storage system or the data written to the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the attribute. 3. The method of any of the preceding statements, wherein the attribute comprises one or more of: a host attribute associated with a host associated with the storage system, the data read from the storage system or the data written to the storage system being associated with the host; an attribute of a source of the read traffic and the write traffic; an attribute of a storage structure within the storage system and from which the data is being read or to which the data is being written; or a storage format attribute associated with a storage format used by the storage system. 4. The method of any of the preceding statements, further comprising: identifying, by the data protection system, a format type of a data instance included in the data written to the storage system; and determining, by the data protection system, that a content of the data instance does not match what would be expected to be received by the storage system for the identified format type; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determination that the content of the data instance does not match what would be expected to be received by the storage system for the identified format type. 5. The method of any of the preceding statements, further comprising: identifying, by the data protection system, a pattern associated with one or more of the read traffic or the write traffic; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the pattern. 6. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the data written to the storage system does not include identifiable header information or that the data written to the storage system includes header information that does not match content included in the data written to the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the data written to the storage system does not include the identifiable header information or that the data written to the storage system includes header information that does not match the content included in the data written to the storage system. 7. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the data read from the storage system is at least partially compressed and includes the identifiable header information; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the data read from the storage system is compressed and includes the identifiable header information. 8. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the data written to the storage system includes data that is not decryptable with a key maintained by an authorized key management system external to the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the data written to the storage system includes data that is not decryptable with the key maintained by the key management system. 9. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the data written to the storage system does not include a correct cryptographic signature

associated with an external data encryption service associated with the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the data written to the storage system does not include the correct cryptographic signature. 10. The method of any of the preceding statements, further comprising: determining, by the data protection system, that data already stored by the storage system is deleted or overwritten by the data written to the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the data already stored by the storage system is deleted or overwritten by the data written to the storage system. 11. The method of any of the preceding statements, further comprising: accessing, by the data protection system, phone home data transmitted by the storage system; and detecting, by the data protection system based on the phone home data, an anomaly associated with the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the detected anomaly. 12. The method of any of the preceding statements, wherein the detecting of the anomaly comprises determining that an overall compressibility of data stored by the storage system is below a historical norm associated with one or more of the storage system or a different storage system. 13. The method of any of the preceding statements, further comprising: detecting, by the data protection system, a rate at which data is read from the storage system and written back to the storage system in encrypted form; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the detected rate. 14. The method of any of the preceding statements, further comprising: inputting, by the data protection system, data representative of one or more attributes of the read traffic, the write traffic, or the storage system into a machine learning model; wherein the determining that the storage system is possibly being targeted by the security threat is further based on an output of the machine learning model. 15. The method of any of the preceding statements, further comprising: determining, by the data protection system, an anomaly in a garbage collection process performed by the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining of the anomaly in the garbage collection process. 16. The method of any of the preceding statements, further comprising: identifying, by the data protection system, an attribute of an additional storage system configured to replicate data stored by the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the attribute of the additional storage system. 17. The method of any of the preceding statements, wherein the threshold represents one or more of a rate, an aggregate amount, or a difference compared to a historical trend associated with the storage system. 18. The method of any of the preceding statements, further comprising performing, by the data protection system in response to the determining that the storage system is possibly being targeted by the security threat, a remedial action with respect to the storage system. 19. The method of any of the preceding statements, wherein the performing of the remedial action comprises directing the storage system to generate a recovery dataset for data stored by the storage system. 20. The method of any of the preceding statements, wherein the performing of the remedial action comprises further comprises directing the storage system to transmit the recovery dataset to a remote storage system for storage by the remote storage system. 21. The method of any of the preceding statements, wherein the transmitting of the recovery dataset to the remote storage system is performed using a network file system (NFS) protocol. (173) 122. The method of any of the preceding statements, wherein the performing of the remedial action comprises notifying the remote storage system of the security threat. 23. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the storage system is actually not being targeted by the security threat; and directing, by the data protection system in response to the determining that the host data is actually not being targeted by the security threat, the storage system to delete the recovery dataset. 24. The method of any of the preceding statements, wherein the recovery dataset comprises a snapshot of a storage structure

within the storage system. 25. The method of any of the preceding statements, further comprising preventing, by the data protection system, the recovery dataset from being deleted until one or more conditions are fulfilled. 26. The method of any of the preceding statements, further comprising directing, by the data protection system, the storage system to generate recovery datasets over time in accordance with a data protection parameter set, the recovery datasets usable to restore data maintained by the storage system to a state corresponding to a selectable point in time. 27. The method of any of the preceding statements, wherein the performing of the remedial action comprises directing, in response to the determining that the storage system is possibly being targeted by the security threat, the storage system to use one or more of the recovery datasets to restore the data maintained by the storage system to a state that corresponds to a point in time that precedes a point in time at which the data protection system determines that the storage system is possibly being targeted by the security threat. 28. The method of any of the preceding statements, wherein the performing of the remedial action further comprises modifying, in response to the determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets. 29. The method of any of the preceding statements, wherein: the data protection parameter set specifies a retention duration for the recovery datasets, the retention duration defining a duration that each recovery dataset is saved before being deleted; and the modifying of the data protection parameter set comprises one or more of increasing the retention duration or suspending the retention duration so that at least some of the recovery datasets are not deleted without a specific instruction provided by a source that manages the storage system. 30. The method of any of the preceding statements, wherein: the data protection parameter set specifies a recovery dataset generation frequency that defines a frequency at which the recovery datasets are generated; and the modifying of the data protection parameter set comprises increasing the recovery dataset generation frequency. 31. The method of any of the preceding statements, wherein: the data protection parameter set specifies a remote storage frequency that defines a frequency at which a subset of recovery datasets in the recovery datasets are transmitted to a remote storage system connected to the storage system by way of a network; and the modifying of the data protection parameter set comprises modifying the remote storage frequency. 32. The method of any of the preceding statements, further comprising: maintaining, by the data protection system, configuration data for the storage system; and determining, by the data protection system, that the storage system is corrupted due to the security threat; wherein the performing of the remedial action comprises using, in response to the determining that the storage system is corrupted, the configuration data to reconstruct a replacement storage system for the storage system. 33. The method of any of the preceding statements, wherein the performing of the remedial action comprises providing a notification of the security threat. 34. The method of any of the preceding statements, wherein the performing of the remedial action comprises restoring, by the data protection system based on one or more recovery datasets generated by the storage system, data stored by the storage system to an uncorrupted state. 35. The method of any of the preceding statements, wherein the one or more recovery datasets comprise one or more of a recovery dataset generated prior to the determining that the storage system is possibly being targeted by the security threat or a recovery dataset generated after the determining that the storage system is possibly being targeted by the security threat. 36. The method of any of the preceding statements, wherein the recovery dataset generated prior to the determining that the storage system is possibly being targeted by the security threat comprises a provisional ransomware recovery structure that can only be deleted or modified in accordance with one or more ransomware recovery parameters. 37. The method of any of the preceding statements, wherein the one or more ransomware recovery parameters specify a number or type of authenticated entities that have to approve a deletion or modification of the provisional ransomware recovery structure before the provisional ransomware recovery structure can be deleted or modified. 38. The method of any of the preceding statements, wherein the one or more ransomware recovery parameters specify a retention duration before

which the provisional ransomware recovery structure can be deleted or modified. 39. The method of any of the preceding statements, wherein the restoring is further based on a version of the data stored by the storage system that resides on a system other than the storage system. 40. The method of any of the preceding statements, wherein: the determining that the storage system is possibly being targeted by the security threat constitutes a first threat detection process; and the method further comprises performing, by the data protection system in response to performing the first threat detection process, a second threat detection process different than the first threat detection process and configured to either confirm that the storage system is possibly being targeted by the security threat with a higher confidence threat detection than the first threat detection process or determine that the storage system is not being targeted by the security threat. 41. The method of any of the preceding statements, wherein the determining that the storage system is possibly being targeted by the security threat comprises determining that there is a potential data corruption in the storage system, and wherein the method further comprises: analyzing, by the data protection system in response to the detecting of the potential data corruption, metrics of the storage system; and determining, by the data protection system based on the analyzing of the metrics of the storage system, a corruption-free recovery point for potential use to recover from the potential data corruption. 42. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the read traffic is within a threshold amount of the write traffic during the time period; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the determining that the read traffic is within the threshold amount of the write traffic during the time period. 43. The method of any of the preceding statements, further comprising: identifying, by the data protection system, an attribute associated with one or more of the data read from the storage system or the data written to the storage system; presenting, by the data protection system within a graphical user interface displayed by a display device, graphical information associated with the attribute; and receiving, by the data protection system by way of the graphical user interface, user input; wherein the determining that the storage system is possibly being targeted by the security threat is further based on the user input. 44. The method of any of the preceding statements, further comprising using, by the data protection system, an unmanipulable clock source internal to the storage system to track the time period. 45. The method of any of the preceding statements, wherein the data protection system is implemented by a controller within the storage system. 46. The method of any of the preceding statements, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network. 47. The method of any of the preceding statements, wherein the determining that the storage system is possibly being targeted by the security threat comprises determining that ransomware is possibly on the storage system (e.g., that a ransomware attack is possibly in progress or operation against the storage system). 48. A system comprising: a memory storing instructions; a processor communicatively coupled to the memory and configured to execute the instructions to: determine that a total amount of read traffic and write traffic processed by a storage system during a time period exceeds a threshold, the read traffic representing data read from the storage system during the time period and the write traffic representing data written to the storage system during the time period; determine that the write traffic is less compressible than the read traffic; and determine, based on the total amount of read traffic and write traffic exceeding the threshold and on the write traffic being less compressible than the read traffic, that the storage system is possibly being targeted by a security threat. 49. The system of statement 48, implementing any of the methods recited in the preceding statements. 50. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to: determine that a total amount of read traffic and write traffic processed by a storage system during a time period exceeds a threshold, the read traffic representing data read from the storage system during the time period and the write traffic representing data written to the storage system during the time period; determine that the write traffic is less compressible than the

read traffic; and determine, based on the total amount of read traffic and write traffic exceeding the threshold and on the write traffic being less compressible than the read traffic, that the storage system is possibly being targeted by a security threat. 51. The non-transitory computer-readable medium of statement 50, implementing any of the methods recited in the preceding statements.

(174) Additional advantages and features of the present disclosure can be further described by the following statements. 1. A method comprising: performing, by a data protection system for a storage system, a first security threat detection process; determining, by the data protection system based on the performing of the first security threat detection process, that the storage system is possibly being targeted by a security threat; and performing, by the data protection system, a second security threat detection process, the second security threat detection process providing higher confidence threat detection than the first security threat detection process. 2. The method of any of the preceding statements, further comprising confirming, by the data protection system based on the performing of the second security threat detection process, that the storage system is possibly being targeted by the security threat. 3. The method of any of the preceding statements, further comprising: performing, by the data protection system based on the determining that the storage system is possibly being targeted by the security threat, a first remedial action with respect to the storage system; and performing, by the data protection system based on the confirming that the storage system is possibly being targeted by the security threat, a second remedial action with respect to the storage system. 4. The method of any of the preceding statements, wherein the first remedial action is different than the second remedial action. 5. The method of any of the preceding statements, wherein the first remedial action or the second remedial action comprises one or more of providing a notification, generating a first recovery dataset, preventing a second recovery dataset from being deleted or modified, modifying a data protection parameter set for a third recovery dataset, or restoring data stored by the storage system to an uncorrupted state. 6. The method of any of the preceding statements, further comprising determining, by the data protection system based on the performing of the second security threat detection process, that the storage system is not being targeted by the security threat. 7. The method of any of the preceding statements, further comprising reverting back, by the data protection system based on the determining that the storage system is not being targeted by the security threat, to performing the first security threat detection process. 8. The method of any of the preceding statements, further comprising performing, by the data protection system based on the determining that the storage system is possibly being targeted by the security threat, a remedial action with respect to the storage system. 9. The method of any of the preceding statements, wherein the performing of the second security threat detection process is performed in response to the determining that the storage system is possibly being targeted by the security threat. 10. The method of any of the preceding statements, wherein the performing of the second security threat detection process is performed in parallel with the performing of the first security threat detection process. 11. The method of any of the preceding statements, wherein the data protection system is implemented by a controller within the storage system. 12. The method of any of the preceding statements, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network. 13. The method of any of the preceding statements, wherein the determining that the storage system is possibly being targeted by the security threat comprises determining that a ransomware attack is possibly in progress against the storage system. 14. A system comprising: a memory storing instructions; a processor communicatively coupled to the memory and configured to execute the instructions to: perform, for a storage system, a first security threat detection process; determine, based on the performing of the first security threat detection process, that the storage system is possibly being targeted by a security threat; and perform a second security threat detection process, the second security threat detection process providing higher confidence threat detection than the first security threat detection process. 15. The system of any of the preceding statements, wherein the processor is further configured to execute the instructions to confirm, based on the performing

of the second security threat detection process, that the storage system is possibly being targeted by the security threat. 16. The system of any of the preceding statements, wherein the processor is further configured to execute the instructions to: perform, based on the determining that the storage system is possibly being targeted by the security threat, a first remedial action with respect to the storage system; and perform, based on the confirming that the storage system is possibly being targeted by the security threat, a second remedial action with respect to the storage system. 17. The system of any of the preceding statements, wherein the first remedial action is different than the second remedial action. 18. The system of any of the preceding statements, wherein the first remedial action or the second remedial action comprises one or more of providing a notification, generating a first recovery dataset, preventing a second recovery dataset from being deleted or modified, modifying a data protection parameter set for a third recovery dataset, or restoring data stored by the storage system to an uncorrupted state. 19. The system of any of the preceding statements, wherein the processor is further configured to execute the instructions to determine, based on the performing of the second security threat detection process, that the storage system is not being targeted by the security threat. 20. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to: perform, for a storage system, a first security threat detection process; determine, based on the performing of the first security threat detection process, that the storage system is possibly being targeted by a security threat; and perform a second security threat detection process, the second security threat detection process providing higher confidence threat detection than the first security threat detection process. (175) Additional advantages and features of the present disclosure can be further described by the following statements. 1. A method comprising: directing, by a data protection system, a storage system to generate recovery datasets over time in accordance with a data protection parameter set, the recovery datasets usable to restore data maintained by the storage system to a state corresponding to a selectable point in time; determining, by the data protection system, that the storage system is possibly being targeted by a security threat; and modifying, by the data protection system in response to the determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets. 2. The method of any of the preceding statements, wherein: the data protection parameter set specifies a retention duration for the recovery datasets, the retention duration defining a duration that each recovery dataset is saved before being deleted; and the modifying of the data protection parameter set comprises one or more of increasing the retention duration or suspending the retention duration so that at least some of the recovery datasets are not deleted without a specific instruction provided by a source that manages the storage system. 3. The method of any of the preceding statements, wherein: the data protection parameter set specifies a recovery dataset generation frequency that defines a frequency at which the recovery datasets are generated; and the modifying of the data protection parameter set comprises increasing the recovery dataset generation frequency. 4. The method of any of the preceding statements, wherein: the data protection parameter set specifies a remote storage frequency that defines a frequency at which a subset of recovery datasets in the recovery datasets are transmitted to a remote storage system connected to the storage system by way of a network; and the modifying of the data protection parameter set comprises modifying the remote storage frequency. 5. The method of any of the preceding statements, further comprising: identifying, by the data protection system, an anomaly with respect to the storage system; wherein the determining that the storage system is possibly being targeted by the security threat is based on the identifying of the anomaly. 6. The method of any of the preceding statements, wherein the identifying of the anomaly comprises: determining that a total amount of read traffic and write traffic processed by the storage system during a time period exceeds a threshold, the read traffic representing data read from the storage system during the time period and the write traffic representing data written to the storage system during the time period; and determining, by the data protection system, that the write traffic is less compressible than the read traffic. 7. The method of

any of the preceding statements, further comprising performing, by the data protection system in response to the determination that the storage system is possibly being targeted by the security threat, an additional remedial action with respect to the storage system. 8. The method of any of the preceding statements, wherein the performing of the additional remedial action comprises directing the storage system to transmit a recovery dataset included in the recovery datasets to a remote storage system for storage by the remote storage system. 9. The method of any of the preceding statements, wherein the performing of the additional remedial action comprises providing a notification of the security threat. 10. The method of any of the preceding statements, wherein the data protection system is implemented by a controller within the storage system. 11. The method of any of the preceding statements, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network. 12. The method of any of the preceding statements, wherein the determining that the storage system is possibly being targeted by the security threat comprises determining that ransomware is possibly on the storage system. 13. The method of any of the preceding statements, further comprising using, by the data protection system, at least one of the recovery datasets to restore the data maintained by the storage system to the state corresponding to the selectable point in time. 14. The method of any of the preceding statements, wherein the determining that the storage system is possibly being targeted by the security threat is performed while the recovery datasets are being generated. 15. A system comprising: a memory storing instructions; a processor communicatively coupled to the memory and configured to execute the instructions to: direct a storage system to generate recovery datasets over time in accordance with a data protection parameter set, the recovery datasets usable to restore data maintained by the storage system to a state corresponding to a selectable point in time; determine that the storage system is possibly being targeted by a security threat; and modify, by in response to the determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets. 16. The system of any of the preceding statements, wherein: the data protection parameter set specifies a retention duration for the recovery datasets, the retention duration defining a duration that each recovery dataset is saved before being deleted; and the modifying of the data protection parameter set comprises one or more of increasing the retention duration or suspending the retention duration so that at least some of the recovery datasets are not deleted without a specific instruction provided by a source that manages the storage system. 17. The system of any of the preceding statements, wherein: the data protection parameter set specifies a recovery dataset generation frequency that defines a frequency at which the recovery datasets are generated; and the modifying of the data protection parameter set comprises increasing the recovery dataset generation frequency. 18. The system of any of the preceding statements, wherein: the data protection parameter set specifies a remote storage frequency that defines a frequency at which a subset of recovery datasets in the recovery datasets are transmitted to a remote storage system connected to the storage system by way of a network; and the modifying of the data protection parameter set comprises modifying the remote storage frequency. 19. The system of any of the preceding statements, wherein: the processor is further configured to execute the instructions to identify an anomaly with respect to the storage system; and the determining that the storage system is possibly being targeted by the security threat is based on the identifying of the anomaly. 20. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to: direct a storage system to generate recovery datasets over time in accordance with a data protection parameter set, the recovery datasets usable to restore data maintained by the storage system to a state corresponding to a selectable point in time; determine that the storage system is possibly being targeted by a security threat; and modify, by in response to the determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets.

(176) Additional advantages and features of the present disclosure can be further described by the

following statements. 1. A method comprising: maintaining, by a data protection system, a metric associated with a similar block detection process performed by a storage system with respect to data processed by the storage system, the metric representative of a measure of similar blocks in the data as detected by the similar block detection process; determining, by the data protection system, that the metric changes by more than a threshold amount; and determining, by the data protection system based on the determining that the metric changes by more than the threshold amount, that the data processed by the storage system is possibly being targeted by a security threat. 2. The method of any of the preceding statements, wherein the determining that the metric changes by more than the threshold amount comprises: determining a first value for the metric, the first value associated with a first time period having a particular characteristic; determining a second value for the metric, the second value associated with a second time period having the particular characteristic; and determining that a difference between the first value and the second value is greater than the threshold amount. 3. The method of any of the preceding statements, wherein the particular characteristic is representative of at least one of a particular day of the week, a particular time period during a day, or a time duration. 4. The method of any of the preceding statements, further comprising setting, by the data protection system, the threshold amount based on one or more attributes associated with the data. 5. The method of any of the preceding statements, further comprising performing, in response to the determination that the data processed by the storage system is possibly being targeted by the security threat, a remedial action with respect to the storage system. 6. The method of any of the preceding statements, wherein the performing the remedial action comprises directing the storage system to generate a recovery dataset for data stored by the storage system. 7. The method of any of the preceding statements, further comprising directing, by the data protection system, the storage system to generate recovery datasets over time in accordance with a data protection parameter set, the recovery datasets usable to restore data maintained by the storage system to a state corresponding to a selectable point in time. 8. The method of any of the preceding statements, wherein the performing the remedial action comprises directing, in response to the determining that the storage system is possibly being targeted by the security threat, the storage system to use one or more of the recovery datasets to restore the data maintained by the storage system to a state that corresponds to a point in time that precedes a point in time at which the data protection system determines that the data processed by the storage system is possibly being targeted by the security threat. 9. The method of any of the preceding statements, wherein the performing the remedial action further comprises modifying, in response to the determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets. 10. The method of any of the preceding statements, wherein: the data protection parameter set specifies a retention duration for the recovery datasets, the retention duration defining a duration that each recovery dataset is saved before being deleted; and the modifying the data protection parameter set comprises one or more of increasing the retention duration or suspending the retention duration so that at least some of the recovery datasets are not deleted without a specific instruction provided by a source that manages the storage system. 11. The method of any of the preceding statements, wherein the performing the remedial action comprises providing a notification of the security threat. 12. The method of any of the preceding statements, further comprising: determining, by the data protection system, that the metric changes by more than an additional threshold amount that is greater than the threshold amount; and performing, by the data protection system based on the determining that the metric changes by more than the additional threshold amount, an additional remedial action with respect to the storage system. 13. The method of any of the preceding statements, wherein the determining that the metric changes by more than the threshold amount comprises determining that the metric differs by more than a threshold amount from a baseline metric associated with a different dataset. 14. A system comprising: a memory storing instructions; and one or more processors communicatively coupled to the memory and configured to execute the instructions to

perform a process comprising: maintaining a metric associated with a similar block detection process performed by a storage system with respect to data processed by the storage system, the metric representative of a measure of similar blocks in the data as detected by the similar block detection process; determining that the metric changes by more than a threshold amount; and determining, based on the determining that the metric changes by more than the threshold amount, that the data processed by the storage system is possibly being targeted by a security threat. 15. The system of any of the preceding statements, wherein the determining that the metric changes by more than the threshold amount comprises: determining a first value for the metric, the first value associated with a first time period having a particular characteristic; determining a second value for the metric, the second value associated with a second time period having the particular characteristic; and determining that a difference between the first value and the second value is greater than the threshold amount. 16. The system of any of the preceding statements, wherein the particular characteristic is representative of at least one of a particular day of the week, a particular time period during a day, or a time duration. 17. The system of any of the preceding statements, wherein the process further comprises setting the threshold amount based on one or more attributes associated with the data. 18. The system of any of the preceding statements, wherein the determining that the metric changes by more than the threshold amount comprises determining that the metric differs by more than a threshold amount from a baseline metric associated with a different dataset. 19. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to perform a process comprising: maintaining a metric associated with a similar block detection process performed by a storage system with respect to data processed by the storage system, the metric representative of a measure of similar blocks in the data as detected by the similar block detection process; determining that the metric changes by more than a threshold amount; and determining, based on the determining that the metric changes by more than the threshold amount, that the data processed by the storage system is possibly being targeted by a security threat. 20. The non-transitory computer-readable medium of any of the preceding statements, wherein the determining that the metric changes by more than the threshold amount comprises: determining a first value for the metric, the first value associated with a first time period having a particular characteristic; determining a second value for the metric, the second value associated with a second time period having the particular characteristic; and determining that a difference between the first value and the second value is greater than the threshold amount. (177) Additional advantages and features of the present disclosure can be further described by the following statements. 1. A method comprising: detecting, by a data protection system, an anomaly associated with a managed run-time environment; determining, by the data protection system based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; and performing, by the data protection system based on the determining that the managed run-time environment is possibly being targeted by the security threat, a remedial action associated with the managed run-time environment, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity. 2. The method of any of the preceding statements, wherein the detecting the anomaly associated with the managed run-time environment comprises: determining a value of a metric associated with the managed run-time environment; and determining that the value of the metric differs by more than a threshold amount from a baseline value of the metric associated with the managed run-time environment. 3. The method of any of the preceding statements, wherein the metric comprises central processing unit (CPU) usage by the managed run-time environment. 4. The method of any of the preceding statements, wherein the metric is representative of storage system access operations performed by the managed run-time environment with respect to the storage system. 5. The method of any of the preceding statements,

wherein the baseline value is representative of a historical average value for the metric over a predetermined period of time for the managed run-time environment. 6. The method of any of the preceding statements, wherein the baseline value is representative of a value of the metric for an additional managed run-time environment that shares at least one attribute with the managed run-time environment. 7. The method of any of the preceding statements, wherein the data associated with the managed run-time environment comprises data accessed by the managed run-time environment. 8. The method of any of the preceding statements, wherein the data associated with the managed run-time environment comprises a backing dataset representative of run-time resources that enable operation of the managed run-time environment. 9. The method of any of the preceding statements, wherein the performing the remedial action further comprises generating an additional recovery dataset. 10. The method of any of the preceding statements, wherein the performing the remedial action further comprises disabling the managed run-time environment. 11. The method of any of the preceding statements, wherein the managed run-time environment comprises at least one of a virtual machine or a containerized application. 12. The method of any of the preceding statements, wherein the data protection system is implemented by a controller within the storage system. 13. The method of any of the preceding statements, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network. 14. A system comprising: a memory storing instructions; and one or more processors communicatively coupled to the memory and configured to execute the instructions to perform a process comprising: detecting an anomaly associated with a managed run-time environment; determining, based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; and performing, based on the determining that the managed run-time environment is possibly being targeted by the security threat, a remedial action associated with the managed run-time environment, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity. 15. The system of any of the preceding statements, wherein the detecting the anomaly associated with the managed run-time environment comprises: determining a value of a metric associated with the managed run-time environment; and determining that the value of the metric differs by more than a threshold amount from a baseline value of the metric associated with the managed run-time environment. 16. The system of any of the preceding statements, wherein the metric comprises central processing unit (CPU) usage by the managed run-time environment. 17. The system of any of the preceding statements, wherein the metric is representative of storage system access operations performed by the managed run-time environment with respect to the storage system. 18. The system of any of the preceding statements, wherein the baseline value is representative of a historical average value for the metric over a predetermined period of time for the managed run-time environment. 19. The system of any of the preceding statements, wherein the baseline value is representative of a value of the metric for an additional managed run-time environment that shares at least one attribute with the managed run-time environment. 20. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to perform a process comprising: detecting an anomaly associated with a managed run-time environment; determining, based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; and performing, based on the determining that the managed run-time environment is possibly being targeted by the security threat, a remedial action associated with the managed run-time environment, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset

that cannot be deleted without a specific instruction provided by an authorized entity.

(178) One or more embodiments may be described herein with the aid of method steps illustrating the performance of specified functions and relationships thereof. The boundaries and sequence of these functional building blocks and method steps have been arbitrarily defined herein for convenience of description. Alternate boundaries and sequences can be defined so long as the specified functions and relationships are appropriately performed. Any such alternate boundaries or sequences are thus within the scope and spirit of the claims. Further, the boundaries of these functional building blocks have been arbitrarily defined for convenience of description. Alternate boundaries could be defined as long as the certain significant functions are appropriately performed. Similarly, flow diagram blocks may also have been arbitrarily defined herein to illustrate certain significant functionality.

(179) To the extent used, the flow diagram block boundaries and sequence could have been defined otherwise and still perform the certain significant functionality. Such alternate definitions of both functional building blocks and flow diagram blocks and sequences are thus within the scope and spirit of the claims. One of average skill in the art will also recognize that the functional building blocks, and other illustrative blocks, modules and components herein, can be implemented as illustrated or by discrete components, application specific integrated circuits, processors executing appropriate software and the like or any combination thereof.

(180) While particular combinations of various functions and features of the one or more embodiments are expressly described herein, other combinations of these features and functions are likewise possible. The present disclosure is not limited by the particular examples disclosed herein and expressly incorporates these other combinations.

(181) Malicious entities (e.g., hackers, malware, and/or other entities) may gain unauthorized access to a storage system, such as any of the storage systems described herein. With such access, the malicious entities may target the storage system with a security threat, such as a ransomware attack, a malware attack, and/or one more other operations configured to destroy, modify, render unusable, or otherwise negatively affect the storage system and/or data maintained by the storage system.

(182) The methods and systems described herein may be configured to detect that a storage system is possibly being targeted by a security threat and to perform various remedial actions in response to detecting that the storage system is possibly being targeted by the security threat.

(183) The methods and systems described herein may additionally or alternatively be configured to detect inadvertent corruption and/or deletion of data stored by a storage system, such as caused by administrative or application errors. For example, the methods and systems described herein may implement monitoring for unexpected behaviors and controls on deletion of certain kinds of data.

(184) Various advantages and benefits may be realized in accordance with the methods and systems described herein. For example, by detecting and performing one or more remedial actions with respect to a security threat targeting a storage system, the methods and systems described herein may minimize or eliminate data corruption, structural damage, and/or performance degradation that may occur as a result of the security threat. Moreover, by implementing a data protection system at the storage level, the methods and systems described herein may provide a last line of defense against security threats, or other forms of corrupting actions, should other data security measures taken at levels above the storage level (e.g., at the client, server, application, or network levels) fail to identify and/or thwart the security threats or other forms of corrupting actions. This may improve the operation of computing devices at both the storage level and at other levels above the storage level.

(185) FIG. 4 illustrates an exemplary data protection system **400** (“system **400**”). As shown, system **400** may include, without limitation, a storage facility **402** and a processing facility **404** selectively and communicatively coupled to one another. Facilities **402** and **404** may each include or be implemented by hardware and/or software components (e.g., processors, memories, communication

interfaces, instructions stored in memory for execution by the processors, etc.). In some examples, facilities **402** and **404** may be distributed between multiple devices and/or multiple locations as may serve a particular implementation.

(186) Storage facility **402** may maintain (e.g., store) executable data used by processing facility **404** to perform any of the operations described herein. For example, storage facility **402** may store instructions **406** that may be executed by processing facility **404** to perform any of the operations described herein. Instructions **406** may be implemented by any suitable application, software, code, and/or other executable data instance. Storage facility **402** may also maintain any data received, generated, managed, used, and/or transmitted by processing facility **404**. Storage facility **402** may additionally maintain any other suitable type of data as may serve a particular implementation.

(187) Processing facility **404** may be configured to perform (e.g., execute instructions **406** stored in storage facility **402** to perform) various processing operations described herein. References herein to operations performed by system **400** may be understood to be performed by processing facility **404**.

(188) FIG. 5 illustrates an exemplary configuration **500** in which a storage system **502** processes read traffic and write traffic. The read traffic represents data read from storage system **502** and the write traffic represents data written to storage system **502**.

(189) Storage system **502** may be implemented by any of the storage systems, devices, and/or components described herein. For example, storage system **502** may be implemented by a local storage system (e.g., a storage system located on-site at a customer's premises) and/or by a remote storage system (e.g., a storage system located in the cloud).

(190) As shown, storage system **502** includes a plurality of storage structures **504** (e.g., storage structures **504-1** through **504-N**) and a controller **506**. Storage structures **504** may each include any logical structure within which data may be stored and/or organized. For example, storage structures **504** may include one or more snapshots, volumes, file systems, object stores, object buckets, key value or relational or other databases, backup datasets, objects that manage a group of volumes, container objects, blocks, etc. In some examples, storage structures **504** are maintained in one or more storage elements (e.g., storage arrays, memories, etc.).

(191) Controller **506** may be configured to control operations of elements included in storage system **502** and may be implemented by any suitable combination of processors, operating systems, and/or other components as described herein. In particular, controller **506** may be configured to produce control data **508** configured to control storage structures **504**. For example, control data **508** may be representative of one or more instructions to create, modify, write to, read from, delete, eradicate, and/or otherwise interact with storage structures **504**.

(192) Read traffic may represent data read from storage system **502** by a source (e.g., a host in communication with storage system **502**), and write traffic may represent data written to storage system **502** by the source. Read and write traffic may occur in response to the source transmitting one or more requests to storage system **502**. These requests may include instructions for controller **506** to perform one or more operations. Such operations may include writing data to a storage structure **504**, reading data from a storage structure **504**, deleting data from a storage structure **504**, overwriting data within a storage structure **504**, and/or deleting a storage structure **504** itself.

(193) In some examples, read traffic, write traffic, and/or one or more requests to interface with storage system **502** may originate from a malicious source and be representative of a ransomware attack on any of the components and/or data within storage system **502** and/or any other malicious operation that destroys, modifies, renders unusable, or otherwise affects any of the components and/or data within storage system **502**.

(194) Accordingly, as described herein, system **400** may, in some examples, be configured to monitor the read and write traffic processed by storage system **502** (e.g., by monitoring one or more requests provided by one or more sources to storage system **502**) to ascertain whether storage system **502** is possibly being targeted by a security threat.

(195) In some examples, system **400** is implemented by storage system **502**. For example, system **400** may be at least partially implemented by controller **506**. Additionally or alternatively, system **400** may be at least partially implemented by one or more computing devices or systems separate from and in communication with storage system **502**.

(196) To illustrate, FIG. **6** shows an exemplary configuration **600** in which a cloud-based monitoring system **602** is communicatively coupled to storage system **502** by way of a network **604**. Cloud-based monitoring system **602** may at least partially implement system **400**.

(197) Network **604** may include the Internet, a wide area network, a local area network, a provider-specific wired or wireless network (e.g., a cable or satellite carrier network or a mobile telephone network), a content delivery network, and/or any other suitable network. Data may flow between storage system **502** and cloud-based data monitoring system **602** using any communication technologies, devices, media, and protocols as may serve a particular implementation.

(198) Cloud-based monitoring system **602** may be implemented by one or more server-side computing devices configured to communicate with storage system **502** by way of network **604**. For example, cloud-based monitoring system **602** may be implemented by one or more servers or other physical computing devices.

(199) Cloud-based monitoring system **602** may be configured to perform one or more remote monitoring operations with respect to storage system **502**. For example, cloud-based monitoring system **602** may be configured to remotely monitor read and write traffic processed by storage system **502** and/or requests processed by storage system **502**. To this end, as shown, cloud-based monitoring system **602** may receive phone-home data **606** from controller **506** of storage system **502** by way of network **604**. Phone-home data **606** may include various types of data that may be used by cloud-based monitoring system **602** to monitor various types of operations performed by storage system **502**. In particular, phone-home data **606** may include data representative of one or more metrics and/or attributes associated with read and write traffic, one or more metrics and/or attributes associated with components within storage system **502**, one or more requests provided by a source to storage system **502**, and/or any other data as may serve a particular implementation.

(200) As shown, cloud-based monitoring system **602** may include a processor **608** configured to process phone-home data **606**. Processor **608** may process phone-home data **606** in any suitable manner. For example, processor **608** may determine, based on phone-home data **606**, that storage system **502** is possibly being targeted by a security threat and transmit instructions **610** to controller **506** to perform one or more remedial actions configured to counteract the security threat.

(201) Various methods that may be performed by system **400** and/or any implementation thereof are described in connection with various flowcharts depicted in the figures. While the flowcharts depicted in the figures illustrate exemplary operations according to one embodiment, other embodiments may omit, add to, reorder, and/or modify any of the operations shown in the flowcharts depicted in the figures. Moreover, each of the operations shown in the flowcharts depicted in the figures may be performed in any of the ways described herein.

(202) FIG. **7** illustrates an exemplary method **700** of dealing with a possible security threat attack against a storage system (e.g., storage system **502**). At operation **702**, system **400** identifies an anomaly associated with a storage system. At operation **704**, system **400** determines, based on the identified anomaly, that the storage system is possibly being targeted by a security threat. At operation **706**, system **400** performs a remedial action (e.g., in response to determining that the storage system is possibly being targeted by the security threat). Examples of each of these operations are described herein.

(203) Various ways in which system **400** may identify an anomaly associated with a storage system and determine that the storage system is possibly being targeted by a security threat are described in connection with FIGS. **8-23**. Each of the processes described in connection with these figures may be performed independently to determine that a storage system is possibly being targeted by a security threat. Alternatively, any number of the processes described in connection with these

figures may be performed concurrently and/or sequentially in any order to determine that a storage system is possibly being targeted by a security threat.

(204) FIG. 8 illustrates an exemplary traffic-based security threat detection method **800** that may be performed by system **400** and/or any implementation thereof. Method **800** may be used alone or in combination with any of the other security threat detection methods described herein.

(205) At operation **802**, system **400** monitors read traffic and write traffic processed by a storage system during a time period. The read traffic represents data read from the storage system during the time period and the write traffic represents data written to the storage system during the same time period. In some examples, system **400** is configured to use an unmanipulable clock source internal to the storage system to track the time period.

(206) System **400** may monitor read and write traffic in any suitable manner. For example, system **400** may analyze metrics generated by the storage system and/or a cloud-based monitoring system (e.g., cloud-based monitoring system **602**) that are representative of an amount of read and write traffic, the type of data included in the read and write traffic, a source of the read and/or write traffic, timestamp data indicative of a date and/or time that the read and write traffic occurs, and/or any other attribute of the read and write traffic as may serve a particular implementation.

(207) The time period during which system **400** monitors the read and write traffic may be of any suitable duration. In some examples, the time period may be set in response to user input (e.g., by an administrator). Additionally or alternatively, the time period may be set and/or adjusted automatically by system **400** based on an occurrence of one or more events and/or based on one or more attributes associated with the read and/or write traffic.

(208) At decision **804**, system **400** determines whether a total amount of read and write traffic exceeds a threshold. At decision **806**, system **400** determines whether the write traffic is less compressible than the read traffic. If the total amount of read and write traffic exceeds the threshold (“Yes” at decision **804**) and the write traffic is less compressible, or has a far high number of incompressible blocks than the read traffic, than the read traffic (“Yes” at decision **806**), system **400** determines at operation **808** that the storage system is possibly being targeted by a security threat. This is because rewriting data as encrypted data is typical of a ransomware attack, and encrypted data is generally not very compressible (in some cases, encrypted data is entirely incompressible). Otherwise, system **400** continues monitoring the read and write traffic (“No” at decision **804** and/or decision **806**).

(209) The threshold to which system **400** compares the total amount of read and write traffic may be any suitable value and type. For example, the threshold may be a particular amount of bytes of data included in the read and write traffic during the time period. Additionally or alternatively, the threshold may be representative of a rate (e.g., a certain amount of data per second, minute, hour, or some other time increment). Additionally or alternatively, the threshold may be representative of aggregate amount (e.g., a total number of bytes). Additionally or alternatively, the threshold may be representative of a difference from historical trends. As an illustration, if system **400** detects a spike in a total amount of read and write traffic during a particular time period compared to a similar time period on a different day, this may be indicative of a possible security threat against the storage system.

(210) In some examples, the threshold to which system **400** compares the total amount of read and write traffic may be set in response to user input (e.g., by an administrator). Additionally or alternatively, the threshold may be set and/or adjusted automatically by system **400** based on an occurrence of one or more events and/or based on one or more attributes associated with the read and/or write traffic. For example, the threshold may be increased during periods of time when the total amount of read and write traffic are typically higher than average. Likewise, the threshold may be decreased during periods of time when the total amount of read and write traffic are typically lower than average.

(211) In some examples, system **400** may maintain data representative of multiple thresholds each

corresponding to different types of data included in the read and write traffic and/or to any other attribute of the read and write traffic. In these examples, system **400** may concurrently compare different segments of the read and write traffic to the different thresholds. If one or more of the thresholds are met, system **400** may satisfy decision **804** (e.g., by proceeding along the “Yes” branch of decision **804**).

(212) System **400** may determine whether the write traffic is less compressible than the read traffic in any suitable manner. For example, system **400** may determine an overall compressibility (e.g., in terms of percentage and/or total amount of storage space saved if compressed) of the write traffic and of the read traffic during the time period. If the overall compressibility of the write traffic is less than the overall compressibility of the read traffic (e.g., by more than a particular threshold), this may indicate that the write traffic includes encrypted data (which has a relatively low amount of compressibility), which may be indicative of a ransomware attack and/or any other type of security threat. It will be recognized that overall compressibility is only one metric that may be used to determine whether the write traffic is less compressible than the read traffic. Other metrics may include file by file comparisons of compressibility, peak compressibility metrics, etc.

(213) As indicated at operation **808**, based on the total amount of read traffic and write traffic exceeding the threshold and on the write traffic being less compressible than the read traffic, system **400** may determine that the storage system is possibly being targeted by a security threat. System **400** may take one or more other factors into consideration when determining whether the storage system is possibly being targeted by the security threat.

(214) For example, FIG. **9** illustrates another exemplary traffic-based security threat detection method **900** that may be performed by system **400** and/or any implementation thereof. Method **900** may be used alone or in combination with any of the other security threat detection methods described herein.

(215) Method **900** is similar to method **800**, except that method **900** further includes another condition that needs to be satisfied before system **400** determines that the storage system is possibly being targeted by the security threat. In particular, at decision **902**, system **400** determines whether the read traffic is within a threshold amount of the write traffic during the time period. This threshold amount may be relatively small such that satisfaction of this condition occurs when the total amount of read traffic during the time period is approximately the same as the total amount of write traffic during the time period. This may be indicative of a ransomware attack or other security threat against the storage system in which data maintained by the storage system is being read out, encrypted, and written back to the storage system.

(216) Hence, if system **400** determines that the read traffic is within the threshold amount of the write traffic during the time period (“Yes” at decision **902**), and if the results of decisions **804** and **806** are both “Yes” as described in connection with FIG. **8**, system **400** may determine that the storage system is possibly being targeted by a security threat.

(217) FIG. **10** illustrates an exemplary attribute-based security threat detection method **1000** that may be performed by system **400** and/or any implementation thereof. Method **1000** may be used alone or in combination with any of the other security threat detection methods described herein. For example, method **1000** may be used in combination with method **800** and/or method **900** to detect data substreams within all of the read and write traffic to a storage system that may be indicative of a methodical attempt by a malicious entity to corrupt data maintained by the storage system (e.g., by encrypting data and/or overwriting a collection of unencrypted data).

(218) At operation **1002**, system **400** identifies an attribute associated with read traffic and/or write traffic processed by a storage system. This identification may be performed while system **400** is monitoring the read and/or write traffic processed by the storage system as described herein. At operation **1004**, system **400** determines, based on the identified attribute, that the storage system is possibly being targeted by a security threat.

(219) The attribute identified in operation **1002** may be any suitable attribute as may serve a

particular implementation. For example, the attribute may include a host attribute that identifies a particular host associated with the storage system. In this implementation, system **400** may monitor host-specific data read from the storage system and/or host-specific data written to the storage system to detect a possible security threat against the storage system. This may be beneficial if there are multiple hosts associated with a particular storage system. In this scenario, host-specific data associated with each host may be monitored in accordance with a different rule set specific to each host.

(220) To illustrate, a particular host may be associated with highly sensitive data (e.g., financial data or other types of personal data) maintained by a storage system that may be more prone to a ransomware attack and/or other type of security threat than other data that is not as sensitive. In this example, a relatively stringent rule set (e.g., a relatively low threshold for decision **804**) may be used when monitoring read and/or write traffic associated with this host. For example, a relatively stringent rule set may be used for a host that does not normally issue traffic to a particular dataset.

(221) As another example, the attribute identified in operation **1002** may include an attribute of a storage structure (e.g., storage structure **504**) within the storage system. For example, the attribute may include an identifier of a particular volume and/or other type of storage structure within the storage system to which data is being written and/or from which data is being read, a storage capacity of a storage structure to which data is being written and/or from which data is being read, and/or any other suitable attribute associated with a particular storage structure.

(222) To illustrate, system **400** may monitor read and write traffic associated with a particular volume within a storage system to determine whether a total amount of read and write traffic exceeds a threshold (decision **804**) and/or whether the write traffic is less compressible than the read traffic (decision **806**). In this manner, a security threat that targets a particular storage structure within a storage system may be more effectively detected.

(223) As another example, the attribute identified at operation **1002** may include a storage format attribute that identifies and/or is otherwise associated with a storage format used by the storage system. For example, the storage format attribute may indicate that the storage system is using an object storage format, a block storage format, and/or a file storage format. This data may be used in any manner to more specifically specify a rule set used to monitor for possible security threats against the storage system.

(224) In some examples, such as in a file-based and/or object-based storage system, stored data (e.g., files and/or objects) may be identifiable as being of a particular type (e.g., an image file, a video file, a ZIP archive, a text file, a machine code binary file, a log file, a database table space, etc.). However, the content of the data may instead look like encrypted data (e.g. randomized and incompressible content) that does not match what would be expected of the particular type. System **400** may be configured to detect these types of content versus format type mismatches and, based on one or more of the mismatches, determine that the storage system is possibly being targeted by a security threat.

(225) To illustrate, FIG. **11** shows an exemplary format type-based security threat detection method **1100** that may be performed by system **400** and/or any implementation thereof. Method **1100** may be used alone or in combination with any of the other security threat detection methods described herein.

(226) At operation **1102**, system **400** monitors write traffic processed by a storage system. This may be performed in any of the ways described herein.

(227) At operation **1104**, system **400** identifies a format type of a data instance (e.g., a file and/or object) included in the write traffic. The format type may be indicative of a particular type of data (e.g., an image file, a database tablespace, etc.). System **400** may identify the format type based on metadata associated with the data instance, a file extension of the data instance, and/or in any other suitable manner.

(228) At decision **1106**, system **400** determines whether the content of the data instance matches

what is expected for the identified format type. If the content of the data instance does not match what is expected for the identified format type (“No” at decision **1106**), system **400** may, at operation **1108**, determine that the storage system is possibly being targeted by a security threat. If the content of the data instance does match what is expected for the identified format type (“Yes” at decision **1106**), system continues monitoring the write traffic processed by the storage system. In some examples, a threshold number of mismatches between data instances and identified format types may be detected before system **400** determines that the storage system is possibly being targeted by a security threat.

(229) FIG. **12** illustrates an exemplary pattern-based security threat detection method **1200** that may be performed by system **400** and/or any implementation thereof. Method **1200** may be used alone or in combination with any of the other security threat detection methods described herein.

(230) At operation **1202**, system **400** identifies a pattern associated with read traffic and/or write traffic processed by a storage system. This identification may be performed while system **400** is monitoring the read and/or write traffic processed by the storage system as described herein. At operation **1204**, system **400** determines, based on the identified pattern, that the storage system is possibly being targeted by a security threat.

(231) System **400** may identify the pattern at operation **1202** in any suitable manner. For example, a ransomware attack may repeatedly read data and write the same data in encrypted form in an identifiable pattern of read/writes. This pattern may be identified by system **400** based on one or more metrics associated with the read and write traffic and used to determine that the storage system is possibly being targeted by a security threat. Such metrics may be included in data maintained by a controller of a storage system and/or in phone home data transmitted by the storage system to a cloud-based monitoring system.

(232) As another example, system **400** may identify a pattern involving reading from a block volume in some pattern with direct overwrites of compressible data with incompressible data (e.g., with no identifiable data format headers, or with incompressible content versus prior content that was compressible) a short time later, particularly in a sequential pattern of reads and a trailing pattern of sequential overwrites. To illustrate, such a pattern may include a read of the first few blocks of a block volume or partition or another recognizable structure stored on a block volume that may be the start of a file system or host-based block device (e.g., a logical volume in a volume manager), followed by an overwrite of that data with relatively incompressible data. In some examples, such a pattern may begin at logical block address (LBA) zero.

(233) As another example, system **400** may identify any pattern of reading unmapped blocks and rewriting of those same blocks with relatively incompressible data, or writing an equivalent amount of relatively incompressible data elsewhere in the storage system.

(234) These patterns, as well as others that may be detected by system **400**, are not common I/O patterns for a storage system and may accordingly be flagged by system **400** as being indicative of a possible security threat against the storage system.

(235) System **400** may detect a pattern indicative of a possible security threat against a storage system over any suitable amount of time. For example, some patterns may be relatively subtle and therefore detected by system **400** over a relatively long amount of time using one or more metrics, machine learning algorithms, and/or other detection algorithms. Other patterns may be detected relatively quickly by system **400**.

(236) In some examples, a confidence level of the determination made by system **400** that the storage system is possibly being targeted by a security threat may change over time as one or more patterns are detected and/or tracked by system **400**. For example, a detected pattern may result in system **400** determining that the storage system is possibly being targeted by a security threat with an initial confidence level. Over time, if the pattern persists or becomes more prevalent, the confidence level of the determination that the storage system is possibly being targeted by the security threat may increase.

(237) FIG. 13 shows an exemplary header information-based security threat detection method **1300** that may be performed by system **400** and/or any implementation thereof. Method **1300** may be used alone or in combination with any of the other security threat detection methods described herein.

(238) At operation **1302**, system **400** monitors read and write traffic processed by a storage system. This may be performed in any of the ways described herein.

(239) At decision **1304**, system **400** determines whether the write traffic includes identifiable header information. If the write traffic does not include identifiable header information (“No” at decision **1304**), system **400** may, at operation **1308**, determine that the storage system is possibly being targeted by a security threat.

(240) If the write traffic does include identifiable header information (“Yes” at decision **1304**), system **400** determines, at decision **1306**, whether the header information matches content included in data written to the storage system. If the header information does not match content included in data written to the storage system, system **400** may, at operation **1308**, determine that the storage system is possibly being targeted by a security threat. Alternatively, if the header information does match content included in data written to the storage system, system **400** may return to monitoring the write traffic at operation **1302**.

(241) As used herein, header information may refer to supplemental data included in (e.g., placed at a beginning of) a block of data being transmitted to the storage system. The header information may identify a format, type, and/or other attribute of data included in a payload portion of the block of data being transmitted to the storage system. Additionally or alternatively, the header information may include a checksum and/or other data that may be used to test for corrupted data.

(242) In some examples, legitimate data (e.g., data not associated with a security threat) being written to a storage system includes identifiable header information that matches content (e.g., payload content) included in the data being written to the storage system. For example, if the identifiable header information of legitimate data indicates that the payload data has a certain format, the payload data should have that format.

(243) However, data associated with a security threat (e.g., a ransomware attack and/or an attempt to write corrupt data to the storage system) may either not have identifiable header information or include identifiable header information that does not match content included in the data being written to the storage system. For example, data being written to a storage system as part of a security threat against the storage system may include header information related to known image, video, sound, or archive files, but payload data included in the data being written to the storage system may not be of any of those types of files. As another example, files may be renamed as part of a re-encryption performed by a malicious entity. For example, if a collection of JPEG files are rewritten into new files with new names, those names may not indicate that they are JPEG files. To detect this, system **400** may determine that a preponderance of files in a directory tree, for example, had been of a particular set of file types by filename pattern, and that those files are being replaced by new files that no longer have that filename pattern.

(244) Accordingly, if data being written to the storage system does not include identifiable header information, system **400** may flag the data as possibly being representative of a security threat against the storage system. Additionally or alternatively, if data written to the storage system includes identifiable header information, but the header information does not match content included in the data written to the storage, system **400** may flag the data as possibly being representative of a security threat against the storage system.

(245) In some examples, system **400** may determine file formats from header data when reading files. This may be performed when files are written out with a name pattern (such as with a .JPG suffix), by recognizing contents of configuration files (such as a database configuration file identifying certain files or block devices as being used as specific parts of a database (logs, tablespaces, etc.)), and/or in any suitable manner. Accordingly, system **400** may detect a change in

filename pattern by detecting when reads have a particular detectable format but writes do not have the same format, or when writes of files with known filename formats (such as .JPG suffixes or the many other suffixes associated with file types) do not result in files with a recognizable format. In response, system **400** may flag data involved in these writes as possibly being representative of a security threat against the storage system.

(246) In some examples, system **400** may base a determination of whether a storage system is being targeted by a security threat by comparing header information included in read traffic with header information included in write traffic. For example, if data read from the storage system is at least partially compressed (e.g., already compressed image, video, or sound files, or even compressed archives) and includes identifiable header information, but no similar identifiable header information can be found in the data being written to the storage system, this may indicate that the read data is being replaced with encrypted data. Hence, system **400** may in this scenario determine that the storage system is possibly being targeted by a security threat.

(247) FIG. **14** shows an exemplary cryptography-based security threat detection method **1400** that may be performed by system **400** and/or any implementation thereof. Method **1400** may be used alone or in combination with any of the other security threat detection methods described herein.

(248) At operation **1402**, system **400** monitors write traffic processed by a storage system. This may be performed in any of the ways described herein.

(249) At decision **1404**, system **400** determines whether data included in the write traffic is encrypted. If the data is not encrypted (“No” at decision **1404**), system **1402** continues monitoring the write traffic (operation **1402**).

(250) However, if the data is encrypted (“Yes” at decision **1404**), system **400** determines at decision **1406** whether the encrypted data is decryptable using a key maintained by an authorized key management system. If the data is decryptable using a key maintained by the authorized key management system (“Yes” at decision **1406**), system **1402** continues monitoring the write traffic (operation **1402**).

(251) However, if the data is not decryptable using a key maintained by the authorized key management system (“No” at decision **1406**), system **400** may determine, at operation **1408**, that the storage system is possibly being targeted by a security threat.

(252) In this example, the authorized key management system may be implemented by any suitable entity and/or system external to and in communication with the storage system. For example, the authorized key management system may utilize the key management interoperability protocol (KMIP) to encrypt legitimate data before the legitimate data is written to the storage system. In some examples, an authorized key management system external to the storage system may facilitate data in motion security before the data is written to the storage system. In some examples, such data in motion security may not prevent system **400** from profiling the underlying data (e.g., the data that has been encrypted) for compressibility.

(253) System **400** may determine whether data included in the write traffic is decryptable into recognizable unencrypted data using a key maintained by an authorized key management system in any suitable manner. For example, system **400** may route the write traffic through the authorized key management system before allowing the write traffic to be written to the storage system. The authorized key management system may determine whether the write traffic is decryptable in any suitable manner. As another example, system **400** may maintain a copy of the key maintained by the authorized key management system and perform any suitable process configured to determine whether the data included in the write traffic is decryptable using the key. As another example, there could be multiple keys that might be used to encrypt data, where the key used for encryption of a particular data item is not obvious from the item itself. Multiple candidate keys could be tried for decryption, as a result, to determine if any of them can decrypt the data into a form that is recognizable as unencrypted.

(254) If the data included in the write traffic is not decryptable by any of several candidate keys

maintained by the authorized key management system, system **400** may determine that the data is possibly associated with a security threat against the storage system.

(255) FIG. **15** shows another exemplary cryptography-based security threat detection method **1500** that may be performed by system **400** and/or any implementation thereof. Method **1500** may be used alone or in combination with any of the other security threat detection methods described herein.

(256) At operation **1502**, system **400** monitors write traffic processed by a storage system. This may be performed in any of the ways described herein.

(257) At decision **1504**, system **400** determines whether data included in the write traffic is encrypted. If the data is not encrypted (“No” at decision **1504**), system **1502** continues monitoring the write traffic (operation **1502**).

(258) However, if the data is encrypted (“Yes” at decision **1504**), system **400** determines at decision **1506** whether the encrypted data includes a correct cryptographic signature. As used herein, a cryptographic signature may refer to any sequence of data (e.g., a digital signature) that indicates that data has been encrypted using a key maintained by an authorized key management system.

(259) If the encrypted data does include a correct cryptographic signature (“Yes” at decision **1506**), system **1502** continues monitoring the write traffic (operation **1502**).

(260) However, if the encrypted data does not include a correct cryptographic signature (“No” at decision **1506**), system **400** may determine, at operation **1508**, that the storage system is possibly being targeted by a security threat.

(261) In some examples, method **1400** and/or method **1500** may be leveraged to provide an end-to-end authentication heuristic from applications through the storage stack to prevent an unauthenticated process from writing data associated with a security threat (e.g., ransomware blocks) to the storage system in the first place.

(262) FIG. **16** shows an exemplary stored data-based security threat detection method **1600** that may be performed by system **400** and/or any implementation thereof. Method **1600** may be used alone or in combination with any of the other security threat detection methods described herein.

(263) At operation **1602**, system **400** monitors write traffic processed by a storage system. This may be performed in any of the ways described herein.

(264) At decision **1604**, system **400** determines whether data already stored by the storage system is being deleted or overwritten by the write traffic. If data is not being deleted or overwritten (“No” at decision **1604**), system **400** continues monitoring the write traffic at operation **1602**.

(265) However, if data is being deleted or overwritten by the write traffic (“Yes” at decision **1604**), system **400** may determine, at operation **1606**, that the storage system is possibly being targeted by a security threat.

(266) System **400** may determine that data already stored by the storage system is being deleted or overwritten by write traffic in any suitable manner. For example, in a file or object based storage system, deletions and overwrites can be detected directly. In the case of an object based storage system that is being used by a host to store file systems or databases, deletions may be inferred by system **400** by a combination of previously read data being overwritten quickly, or sometime later (such as because blocks added to a free list were eventually reused), by being unmapped, or by being overwritten with zeros. Such deletions or overwrites may in and of themselves be indicative of a possible security threat against a storage system. Additionally or alternatively, such deletions or overwrites in combination with any of the other security threat detection methods described herein may be indicative of a possible security threat against a storage system.

(267) FIG. **17** shows a remote security threat detection method **1700** that may be performed by system **400** and/or any implementation thereof. Method **1700** may be used alone or in combination with any of the other security threat detection methods described herein.

(268) At operation **1702**, system **400** accesses phone home data (e.g., phone home data **606**) transmitted by a storage system. This may be performed in any of the ways described herein.

(269) At operation **1704**, system **400** detects, based on the phone home data, an anomaly associated with the storage system. The anomaly may include any of the anomalies described herein.

(270) At operation **1706**, system **400** determines, based on the detected anomaly, that the storage system is possibly being targeted by a security threat.

(271) To illustrate, a cloud-based monitoring system implementation of system **400** may use the phone home data transmitted thereto by a storage system to identify a pattern and or attribute of read and/or write traffic that may be indicative of a possible security threat against the storage system. For example, system **400** may detect that an overall compressibility of data stored by the storage system is below a historical norm associated with the storage system or with a different storage system (e.g., a different storage system that has one or more similar attributes as the storage system). Based on this, system **400** may determine that the storage system is possibly being targeted by a security threat.

(272) In some examples, system **400** may be provided with user input that identifies certain metrics that system **400** should focus on (e.g., in phone home data and/or in metrics data maintained by the storage system) when monitoring for anomalies that may be indicative of a security threat against the storage system. For example, a customer of a storage system may provide user input representative of expected types of data for the write traffic so that system **400** may take that information into account when analyzing the write traffic.

(273) FIG. **18** shows an exemplary rate-based security threat detection method **1800** that may be performed by system **400** and/or any implementation thereof. Method **1800** may be used alone or in combination with any of the other security threat detection methods described herein.

(274) At operation **1802**, system **400** detects a rate at which data is read from a storage system and written back to the storage system in encrypted form. At operation **1804**, system **400** determines, based on the detected rate, that the storage system is possibly being targeted by a security threat.

(275) To illustrate, some relatively slow write patterns may be an indication that a malicious entity is in fact doing something that a normal program that encrypts a dataset for legitimate purposes would not be doing. For example, a sequential process that reads data and writes back that same data in incompressible form, but that does so at a rate that is much slower than a legitimate process would do so may itself be an indication that the storage system is being targeted by a security threat. System **400** may be configured to detect this difference in rate and, in response, determine that the storage system is possibly being targeted by a security threat.

(276) As another example, a process that slowly rewrites a set of files that were originally in a recognizable format into a nonrecognizable format may be an indication that the storage system is being targeted by a security threat. Based on this relatively slow rewrite process, system **400** may determine that the storage system is possibly being targeted by a security threat.

(277) As another example, a read/write process that is relatively faster than what would be expected during a particular time period may be an indication that the storage system is being targeted by a security threat. For example, during a weekend when read/write traffic is historically relatively slow, if system **400** detects a rate of read/writes that is above a particular threshold, system **400** may determine that the storage system is possibly being targeted by a security threat.

(278) Rate-based detection of security threats may be performed over any suitable amount of time. For example, to detect relatively slow rates, system **400** may monitor one or more metrics associated with read/write traffic over the course of a relatively long period of time. In these cases, system **400** may lock down and/or otherwise maintain one or more recovery datasets (e.g., provisional ransomware recovery structures, as described herein) for a relatively long period of time in case they are needed to recover from data corruption caused by the security threat.

(279) FIG. **19** shows an exemplary machine learning model-based security threat detection method **1900** that may be performed by system **400** and/or any implementation thereof. Method **1900** may be used alone or in combination with any of the other security threat detection methods described herein.

(280) At operation **1902**, a machine learning model is trained (e.g., by system **400** and/or any other system) to detect anomalies associated with read/write traffic processed by a storage system. The machine learning model may be supervised and/or unsupervised as may serve a particular implementation and may be configured to implement one or more decision tree learning algorithms, association rule learning algorithms, artificial neural network learning algorithms, deep learning algorithms, bitmap algorithms, and/or any other suitable data analysis technique as may serve a particular implementation. In some examples, the machine learning model is trained with actual ransomware payloads.

(281) In some examples, the machine learning model is trained using honeypot files, sectors of blocks, and/or any other data structure configured to serve as a decoy for ransomware and other security threats. These honeypot data structures may be maintained by system **400** at any suitable location (e.g., within the storage system or remote from the storage system). Based on how attackers interact with the honeypot data structures, system **400** may train the machine learning model. The honeypot outputs may additionally or alternatively be used in combination with any of the other security threat detection methods described herein.

(282) At operation **1904**, system **400** inputs attribute data for read traffic, write traffic, and/or the storage system into the machine learning model. The machine learning model may process this attribute data in any suitable manner. For example, the machine learning model may be trained to look at deduplication checksum/hashes in a data reducing storage array leveraging an out of band cloud service that cannot be compromised. This may allow the machine learning model to recognize when write traffic differs from a historical trend. As another example, the machine learning model may be configured to detect actual ransomware payloads within write traffic.

(283) At operation **1906**, system **400** determines, based on an output of the machine learning model, that the storage system is possibly being targeted by a security threat. This may be performed in any suitable manner. For example, the output of the machine learning may include a confidence score. If the confidence score is above a certain threshold, system **400** may determine that the storage system is possibly being targeted by a security threat.

(284) FIG. **20** shows an exemplary garbage collection-based security threat detection method **2000** that may be performed by system **400** and/or any implementation thereof. Method **2000** may be used alone or in combination with any of the other security threat detection methods described herein.

(285) At operation **2002**, system **400** monitors a garbage collection process performed by a storage system. The garbage collection process may include any process configured to reclaim storage system space as described herein.

(286) At decision **2004**, system **400** determines whether there is an anomaly in the garbage collection process performed by the storage system. If system **400** does not determine that there is an anomaly in the garbage collection process performed by the storage system (“No” at decision **2004**), system **400** continues monitoring the garbage collection process at operation **2002**.

However, if system **400** determines that there is an anomaly in the garbage collection process performed by the storage system (“Yes” at decision **2004**), system **400** may determine at operation **2006** that the storage system is possibly being targeted by a security threat.

(287) System **400** may detect an anomaly in a garbage collection process performed by a storage system in any suitable manner. For example, as data in a segment becomes invalid due to a ransomware attack or other security threat, the data becomes more attractive for garbage collection. This may result in a higher than average amount of garbage collection performed by the storage system. This change in garbage collection may be detected by system **400** by analyzing metrics associated with the garbage collection process and may be used to determine that the storage system is possibly being targeted by a security threat.

(288) System **400** may additionally or alternatively monitor one or more other internal processes performed by a storage system to determine whether the storage system is possibly being targeted

by a security threat. For example, system **400** may monitor a deep compression process performed by one or more libraries within a storage system. In this example, system **400** may monitor for block patterns that match one or more rootkits and/or other structures that are put in place during the full lifecycle of a malicious attack. Such block patterns may be indicative of an impending encryption process that happens at the end of the malicious attack. If such block patterns are detected, system **400** may determine that the storage system is possibly being targeted by a security threat.

(289) In some configurations, first and second storage systems are configured to serve as replicating storage systems one for another. For example, any data stored in the first storage system may be replicated in the second storage system. This may provide various data redundancy and security features. In these configurations, system **400** may be configured to identify attributes of both storage systems (e.g., by monitoring read/write traffic at both storage systems) to determine whether one or both of the storage systems are possibly being targeted by a security threat.

(290) To illustrate, FIG. **21** shows an exemplary replicating storage system-based security threat detection method **2100** that may be performed by system **400** and/or any implementation thereof. Method **2100** may be used alone or in combination with any of the other security threat detection methods described herein.

(291) At operation **2102**, system **400** monitors read and/or write traffic processed by a storage system. This may be performed in any of the ways described herein.

(292) At decision **2104**, system **400** determines whether there is an anomaly associated with the storage system (e.g., with the read and/or write traffic processed by the storage system). This may be performed in any of the ways described herein. If system **400** does not detect an anomaly (“No” at decision **2104**), system **400** continues monitoring the read and/or write traffic at operation **2102**.

(293) However, if system **400** detects an anomaly (“Yes” at decision **2104**), system **400** may determine whether a similar (e.g., the same) anomaly exists at a replicating storage system configured to replicate data stored by the storage system (decision **2106**).

(294) If system **400** does not detect an anomaly at the replicating storage system (“No” at decision **2106**), system **400** continues monitoring the read and/or write traffic at operation **2102**.

(295) However, if system **400** detects an anomaly at the replicating storage system (“Yes” at decision **2106**), system **400** may determine at operation **2108**, based on the anomaly being detected at both storage systems, that the storage system (and, in some cases, the replicating storage system) is possibly being targeted by a security threat.

(296) By way of example, analyzers implementing system **400** may be run on both storage systems when a dataset is replicated between a first and second storage system, with results compared so that the two storage systems can serve as checks on each other. For example, metrics that lead system **400** to provisionally determine that the first storage system is possibly being targeted by a security threat may be exchanged to ensure that both storage systems are seeing the same information.

(297) Since many read and write requests (or file system, database, or object requests) may only be received by one of the storage systems, the other storage system may only have some metrics. For example, the second storage system, such as one that is the target of an asymmetric form of replication for a dataset, may only have information on general compressibility of writes for that dataset. The two storage systems may still exchange the metrics they do have with each other, with each comparing the metrics they do know to those coming from the other storage system and using the metrics received from the other storage system.

(298) For example, a combination of actual profiles of read, writes, overwrites, compressibility of data in actual read and write requests, and metrics organized by hosts may be exchanged with the storage systems comparing general compressibility of written data for anomalies between the two storage systems and with the additional information from a first storage system used to duplicate the first storage system's analysis on the second storage system.

(299) In a symmetrically replicated storage system with symmetric access to replicated datasets, metrics may be exchanged to ensure that both storage systems have the relevant data necessary for either of them to detect some types of anomalies, such as because some read, write, or other requests are directed to one of the two storage systems while other read, write, or other requests are directed to the other of the two storage systems.

(300) These kinds of exchanges may further be used to detect some examples where one or the other storage system has been compromised. For example, secure hashes of metrics that both storage systems are expected to know may be exchanged rather than exchanging those metrics directly, so a compromised storage system cannot use trends in a common metric received from the other storage system to guess future values for that metric to fool the paired storage system. Since metrics can have some natural differences between two storage systems such as due to other activity, differences in snapshots, or delays in when updates are received and processed, securely hashed metrics may allow for approximations. This may be done in several ways, such as by providing a small set of secure hashes corresponding to discrete ranges of values. For example, if a compressibility factor or compressibility factors within time ranges is provided for recent updates to a dataset, such as based on a percentage, if a first storage system sees an overall compressibility in recent updates of 20% on 100 MB of updates in the prior 30 second interval, and another sees a compressibility of recent updates of 18% on 99 MB of updates in the prior 30 second interval, then the first storage system may securely hash values representing two compressibility ranges of 18% to 20% and 20% to 22% each combined with two update quantity ranges of 98 MB to 100 MB and 100 MB to 102 MB, (forming four secure hashes of each compressibility range with each update quantity range) and the second storage system may securely hash values representing two compressibility ranges of 16% to 18% and 18% to 20% each combined with three update quantity ranges of 96 MB to 98 MB, 98 MB to 100 MB, and 100 MB to 102 MB (forming six secure hashes of each compressibility range with each update quantity range). Since one of the ranges from the first storage system agrees with one of the ranges from the second storage system, the storage systems can be seen as agreeing closely enough without having exchanged too much data about their actual metrics.

(301) In some examples, metrics may be shared to some third system or to a cloud service or some vendor provided service for comparison purposes, in addition to or rather than the two storage systems themselves sharing these anomaly detection metrics data between them. If the two storage systems do not exchange these metrics, then an external system or service can be more certain that the metrics it is receiving from each system are not being guessed at by compromised storage system based on data it is exchanging with an uncompromised storage system.

(302) FIG. 22 shows an exemplary user input-based security threat detection method **2200** that may be performed by system **400** and/or any implementation thereof. Method **2200** may be used alone or in combination with any of the other security threat detection methods described herein.

(303) At operation **2202**, system **400** identifies an attribute associated with data read from the storage system and/or data written to the storage system. The attribute may include one or more of the attributes described herein.

(304) At operation **2204**, system **400** presents, within a graphical user interface displayed by a display device, graphical information associated with the attribute. For example, system **400** may present one or more graphs, analytics information, etc. associated with the attribute.

(305) At operation **2206**, system **400** receives user input by way of the graphical user interface (and/or by way of any other means for receiving user input, such as by way of an API). For example, a user (e.g., an administrator) may, based on the graphical information, provide user input indicating that the attribute is indicative of a possible security threat against the storage system.

(306) At operation **2208**, system **400** determines, based on the user input, that the storage system is possibly being targeted by a security threat.

(307) To illustrate, system **400** may provide a graph over time of various metrics that may be useful

for determining when an attack may have started over time. Based on this graph, a user may provide user input indicating that the storage system has been possibly targeted by a security threat. (308) In some examples, system **400** may be configured to perform multiple security threat detection processes to determine whether a storage system is being targeted by a security threat. For example, system **400** may perform two or more of the security threat detection methods described in connection with FIGS. **8-22**. These threat detection processes may be performed in parallel and/or serially as may serve a particular implementation.

(309) Some security threat detection processes provide higher confidence threat detection than others. In other words, some security threat detection processes may detect a possible security threat with higher accuracy than others. However, a relatively high confidence threat detection method may, in some instances, be more resource intensive and/or take more time than relatively low confidence threat detection methods. Hence, in some examples, system **400** may be configured to initially use a first security threat detection process to provisionally determine that a storage system is a target of a security threat. System **400** may then use a second security threat detection process that provides higher confidence threat detection than the first security threat detection process to verify the provisional determination that the storage system is a target of the security threat.

(310) To illustrate, FIG. **23** shows an exemplary multi-level security threat detection method **2300** that may be performed by system **400** and/or any implementation thereof. Method **2300** may be used alone or in combination with any of the other security threat detection methods described herein.

(311) At operation **2302**, system **400** performs a first security threat detection process with respect to a storage system. The first security threat detection process may include any of the security threat detection processes described herein.

(312) At decision **2304**, system **400** determines, based on the first security threat detection process, whether the storage system is a possible target of a security threat. If system **400** determines that the storage system is not a possible target of security threat based on the first security threat detection process (“No” at decision **2304**), system **400** continues to perform the first security threat detection process at operation **2302**.

(313) However, if system **400** determines, based on the first security threat detection process, that the storage system is a possible target of a security threat (“Yes” at decision **2304**), system **400** may perform a second security threat detection process with respect to the storage system (operation **2306**). The second security threat detection process is configured to provide higher confidence threat detection than the first security threat detection process. Based on the results of the second security threat detection process, system **400** may either confirm that the storage system is being targeted by the security threat or determine that the storage system is not being targeted by the security threat.

(314) In some examples, the second security threat detection process is performed in response to determining that the storage system is possibly being targeted by the security threat. Alternatively, the second security threat detection process may be performed in parallel with the first second security threat detection process.

(315) In method **2300**, the first and second security threat detection processes may be different in some examples. For example, the first security threat detection process may require less resources to perform than the second security threat detection process. In alternative examples, the first and second security threat detection processes are similar processes. In these examples, the second security threat detection process may, for example, be performed for a longer duration and/or with different parameters to provide the higher confidence threat detection. In alternative examples, the first and second security threat detection processes are the same, just performed over different time periods to make a determination with different levels of accuracy.

(316) Various remedial actions that may be performed by system **400** in response to determining

that a storage system is possibly being targeted by a security threat are described in connection with FIGS. 24-30. Each of the processes described in connection with these figures may be performed independently or in combination (e.g., sequentially or concurrently) with other processes used to perform a remedial action. Moreover, each of the remedial action processes described in connection with these figures may be performed in connection with one or more of the security threat detection processes described herein.

(317) FIG. 24 shows an exemplary recovery dataset-based remedial action method 2400 that may be performed by system 400 and/or any implementation thereof. Method 2400 may be used alone or in combination with any of the other remedial action methods described herein.

(318) At operation 2402, system 400 determines that a storage system is possibly being targeted by a security threat. This may be performed in any of the ways described herein.

(319) At operation 2404, system 400 directs the storage system to generate, in response to the determination that the storage system is possibly being targeted by the security threat, a recovery dataset for data stored by the storage system. The recovery dataset may include a snapshot, a backup dataset, an ordered log of metadata describing an ordered application of updates to data maintained by the storage system, and/or any other suitable data structure that may be used to restore data to an uncorrupted state. The recovery dataset 400 may be for all data stored by the storage system, data stored on a particular storage structure (e.g., a volume), data associated with a particular host, and/or any other subset of data stored by the storage system.

(320) By directing the storage system to immediately generate a recovery dataset in response to determining that the storage system is possibly being targeted by the security threat, system 400 may use the recovery dataset (or direct the storage system to use the recovery dataset) to restore at least some of the data maintained by the storage system to an uncorrupted state should the possible security threat turn out to be an actual security threat. In some examples, the recovery dataset is used in combination with one or more previously generated recovery datasets and/or other data sources (e.g., data residing at a host) to restore data that is already corrupted before the recovery dataset is generated in response to the determination that the storage system is possibly being targeted by the security threat.

(321) In some examples, system 400 may direct the storage system to transmit the recovery dataset to a remote storage system for storage by the remote storage system. The remote storage system may include any combination of computing devices remote from and communicatively coupled to the storage system (e.g., by way of a network). In this manner, the recovery dataset itself may be protected from the security threat. In some examples, the transmission of the recovery dataset to the remote storage system is performed using a NFS protocol, an object store protocol, an SMB storage protocol, an S3 storage protocol, and/or any other storage protocol as may serve a particular implementation. In some examples, the remote storage system is implemented by write-only media with restrictions on deletions.

(322) In some examples, system 400 may notify the remote storage system of the security threat so that the remote storage system may abstain from deleting the recovery dataset until one or more conditions are fulfilled. Such conditions may include, but are not limited to, input provided by one or more authenticated entities, a notification from system 400 that it is safe to delete the recovery dataset, etc. For example, system 400 may determine that the storage system is actually not being targeted by the security threat. In response, system 400 may transmit a command to the remote storage system for the remote storage system to delete the recovery dataset.

(323) In embodiments where the recovery dataset is stored within the storage system, system 400 may prevent the recovery dataset from being deleted or modified until system 400 determines that the recovery dataset is not needed to restore data within the storage system. For example, system 400 may direct the storage system to lock down the recovery dataset, make the recovery dataset read-only, make the recovery dataset hidden, and/or otherwise protect the recovery dataset. When one or more conditions are fulfilled (e.g., input from one or more authenticated entities, passage of

a set amount of time, etc.), system **400** may allow the storage system to delete the recovery dataset. (324) FIG. **25** shows an exemplary continuous data protection-based remedial action method **2500** that may be performed by system **400** and/or any implementation thereof. Method **2500** may be used alone or in combination with any of the other remedial action methods described herein. (325) At operation **2502**, system **400** directs a storage system to generate recovery datasets over time (e.g., as a rolling set of snapshots) in accordance with a data protection parameter set. As described herein, these recovery datasets are usable to restore data maintained by the storage system to a state corresponding to a selectable point in time. The data protection parameter set may define one or more parameters associated with the generation of the recovery datasets over time, as described herein.

(326) At operation **2504**, system **400** determines that the storage system is possibly being targeted by a security threat. This may be performed in any of the ways described herein. In some examples, one or more of the recovery datasets generated at operation **2502** are generated prior to system **400** determining that the storage system is possibly being targeted by the security threat. One or more of the recovery datasets generated at operation **2502** may also be generated subsequent to system **400** determining that the storage system is possibly being targeted by the security threat.

(327) At operation **2506**, system **400** modifies, in response to determining that the storage system is possibly being targeted by the security threat, the data protection parameter set for one or more of the recovery datasets.

(328) To illustrate, the data protection parameter set may specify a retention duration for one or more of the recovery datasets. The retention duration defines a duration that each recovery dataset is saved before being deleted (e.g., 24 or 48 hours, or longer in the case of, for example, a weekend or extended break). In the absence of a detected security threat, each recovery dataset may be retained for only a relatively short duration before being deleted. However, based on a determination that the storage system is possibly being targeted by a security threat, system **400** may either increase the retention duration or suspend the retention duration so that at least some of the recovery datasets are not deleted without a specific instruction provided by a source that manages the storage system. In this manner, one or more of the recovery datasets may be used to restore data on the storage system to an uncorrupted state if system **400** determines that the storage system has in actuality been targeted by the security threat.

(329) As another example, the data protection parameter set may additionally or alternatively specify a recovery dataset generation frequency that defines a frequency at which the recovery datasets are generated. In this example, based on a determination that the storage system is possibly being targeted by a security threat, system **400** may increase the recovery dataset generation frequency so that more recovery datasets are available for use in restoring data on the storage system to an uncorrupted state if system **400** determines that the storage system has in actuality been targeted by the security threat.

(330) As another example, the data protection parameter set may additionally or alternatively specify a remote storage frequency that defines a frequency at which a subset of recovery datasets in the recovery datasets are transmitted to a remote storage system connected to the storage system by way of a network (e.g., by using a network file system, streaming backup, or object storage protocol). In this example, based on a determination that the storage system is possibly being targeted by a security threat, system **400** may modify the remote storage frequency. For example, system **400** may increase the remote storage frequency so that more recovery datasets are stored in a read-only format on the remote storage system and available for use in restoring data on the storage system to an uncorrupted state if system **400** determines that the storage system has in actuality been targeted by the security threat.

(331) System **400** may additionally or alternatively direct the storage system to generate (e.g., periodically and/or in response to an occurrence of certain events) one or more provisional ransomware recovery structures (e.g., snapshots). These provisional ransomware recovery

structures may be configured such that they can only be deleted or modified in accordance with one or more ransomware recovery parameters. For example, the one or more ransomware recovery parameters may specify a number or a collection of types of authenticated entities that have to approve a deletion or modification of a provisional ransomware recovery structure before the provisional ransomware recovery structure can be deleted or modified. As another example, the one or more ransomware recovery parameters may specify a minimum retention duration before which the provisional ransomware recovery structure can be deleted or modified.

(332) In some examples, any of the recovery datasets generated herein may be converted to a set of locked-down snapshots (or other suitable types of recovery datasets), possibly as a combination of discretionary snapshots formed early in a possible attack which can be deleted by the storage system itself if system **400** determines that the detection was a false alarm that did not stand up to deeper scrutiny. Additionally or alternatively, instead of formalizing formation and holds on discretionary snapshots, garbage collection, merges, deletions, and/or other maintenance on continuous data protection stores or frequent snapshots may be put on hold pending further analysis. For example, as soon as a bump in incompressible writes is received that is beyond historical norms, system **400** may initiate discretionary lockdowns to avoid deleting recent recoverable images of the storage system that precede the increase in incompressible writes. If the increase in incompressible writes reverts to the historical norm or does not hold up as sufficient to indicate a plausible ransomware attack, then the discretionary snapshots or holds on maintenance operations may be released. If they do hold up, they may be converted into ransomware/corruption protection snapshots with their increased scrutiny required for deletion (or with no means of deleting them within some designated or scheduled time frame). In the case of a continuous data protection store, rather than forming a ransomware/corruption protection snapshot, cleanup or merger of consistency points within the continuous data protection store itself may be blocked from occurring or severely reduced if a plausible sustained attack is detected, with the same kinds of models for duration of time and change authorization models that are described for ransomware/corruption protection snapshots.

(333) A continuous data protection store is a feature of a storage system that records updates to a dataset in such a way that consistent images of prior contents of the dataset can be accessed with a low time granularity (often on the order of seconds, or even less), and stretching back for a reasonable period of time (often hours or days). These allow access to very recent consistent points in time for the dataset, and also allow access to access to points in time for a dataset that might have just preceded some event that, for example, caused parts of the dataset to be corrupted or otherwise lost, while retaining close to the maximum number of updates that preceded that event. Conceptually, they are like a sequence of snapshots of a dataset taken very frequently and kept for a long period of time, though continuous data protection stores are often implemented quite differently from snapshots. A storage system implementing a data continuous data protection store may further provide a means of accessing these points in time, accessing one or more of these points in time as snapshots or as cloned copies, or reverting the dataset back to one of those recorded points in time.

(334) Over time, to reduce overhead, some points in the time held in a continuous data protection store can be merged with other nearby points in time, essentially deleting some of these points in time from the store. This can reduce the capacity needed to store updates. It may also be possible to convert a limited number of these points in time into longer duration snapshots. For example, such a store might keep a low granularity sequence of points in time stretching back a few hours from the present, with some points in time merged or deleted to reduce overhead for up to an additional day. Stretching back in the past further than that, some of these points in time could be converted to snapshots representing consistent point-in-time images from only every few hours.

(335) FIG. 26 shows an exemplary data restoration method **2600** that may be performed by system **400** and/or any implementation thereof. Method **2600** may be used alone or in combination with

any of the other remedial action methods described herein.

(336) At operation **2602**, system **400** determines that a storage system has been targeted by a security threat. As part of this, system **400** may identify data on the storage that has been corrupted by the security threat.

(337) At operation **2604**, system **400** restores (e.g., by directing the storage system to restore), based on one or more recovery datasets generated by the storage system, data stored by the storage system to an uncorrupted state.

(338) The one or more recovery datasets used to restore the data to the uncorrupted state at operation **2604** may include one or more recovery datasets generated prior to system **400** determining that the storage system is possibly being targeted by the security threat (e.g., a recovery dataset generated in accordance with continuous data protection-based remedial action method **2500** and/or a provisional ransomware recovery structure). Additionally or alternatively, the one or more recovery datasets used to restore the data to the uncorrupted state at operation **2604** may include one or more recovery datasets generated after system **400** determines that the storage system is possibly being targeted by the security threat (e.g., a recovery dataset generated in accordance with recovery dataset-based remedial action method **2400**).

(339) In some examples, system **400** may further perform the data restoration based on a version of the data that resides on a system other than the storage system. This other system may include a replicating storage system, a host computing device, and/or any other suitable system as may serve a particular implementation. For example, host data residing on a host computing device may be used in combination with one or more of the recovery datasets described herein to restore data residing on the storage system to an uncorrupted state.

(340) In some examples, system **400** may select a recovery dataset for use in restoring data to the storage system by first determining a corruption-free recovery point for the storage system. This corruption-free recovery point corresponds to a point in time that precedes any data corruption caused by the security threat. System **400** may then select a recovery dataset that corresponds to the corruption-free recovery point for use in a data restoration process.

(341) System **400** may determine a corruption-free recovery point for a storage system in any suitable manner. For example, FIG. 27 shows an exemplary data restoration method **2700** that may be performed by system **400** and/or any implementation thereof. Method **2700** may be used alone or in combination with any of the other remedial action methods described herein.

(342) At operation **2702**, system **400** detects a potential data corruption in a storage system. The potential data corruption may be caused by any of the security threats described herein. System **400** may detect the potential data corruption based on one or more metrics maintained or generated by the storage system, an analysis of the data stored by the storage system, and one or more attributes of a security threat that causes the potential data corruption.

(343) At operation **2704**, system **400** analyzes, in response to detecting the potential data corruption, one or more metrics of the storage system. These metrics may be any of the metrics described herein.

(344) At operation **2706**, system **400** determines, based on the analyzing of the one or more metrics of the storage system, a corruption-free recovery point for potential use to recover from the potential data corruption. The corruption-free recovery point may be determined automatically by system **400** based on one or more metrics associated with the storage system and/or data maintained by the storage system. Additionally or alternatively, system **400** may determine the corruption-free recovery point based on user input provided by a user.

(345) To illustrate, in some examples, system **400** may present (e.g., within a graphical user interface) one or more visualizations that may assist a user in identifying a corruption-free recovery point. For example, system **400** may visualize changes and/or types of changes either in a continuous data protection store or in a time-ordered set of snapshots. For example, system **400** may provide a graph over time of various metrics that may be useful for determining when an

attack may have started over time, where changes are related to differences between snapshots or between continuous data protection consistency points, presented in time order. These metrics may include reads, writes, compressibility of reads, compressibility of writes, and hosts issuing requests, including possibly a visualization of unusual compressibility ratios for particular datasets and from particular hosts.

(346) If file, object, or database information is available for a set of changes (such as but not exclusively because the storage system is itself a file, database, or object server, or because a storage system hosting a block volume used by a client host to store a file system, object store, or database has suitable format analyzers that can determine file, object/bucket, or database changes from the stored file system or database), system **400** can further add metrics related to number of files or objects or database elements or blobs read, with their compressibility, including from various hosts, and number of files or objects written, overwritten, or created and then written, again including compressibility. A visualizer may provide the ability to zoom into directories, buckets, files, tablespaces, or objects that show activity of interest, and then provide graphs or other visualizations to show activity against those over time, including the ability to segregate by hosts or networks from which requests were received.

(347) By being able to hone in on a particular update stream which seems to be the source of deliberate corruption or encryption, these visualizations can be used by a user to trace back in time to when that activity may have started. Then, a continuous data protection consistency point or snapshot from prior to that can be used as a corruption-free starting point for recovery from the attack. Further, system **400** can visualize activity from the hosts used for the attack to isolate which parts of the storage system may have been attacked and corrupted or encrypted, which should suggest that other stored data was not affected and can likely be considered safe.

(348) FIG. **28** shows an exemplary replacement storage system reconstruction method **2800** that may be performed by system **400** and/or any implementation thereof. Method **2800** may be used alone or in combination with any of the other remedial action methods described herein.

(349) At operation **2802**, system **400** maintains configuration data for a storage system. The configuration data may include data representative of one or more host connections and identities, storage system target endpoint addresses, and/or other types of configuration information for a storage system.

(350) At operation **2804**, system **400** determines that the storage system is corrupted due to a security threat. This determination may be performed in any suitable manner.

(351) At operation **2806**, system **400** uses the configuration data to reconstruct a replacement storage system for the storage system. The replacement storage system may be separate from the storage system and/or within the same storage system as may serve a particular implementation.

(352) FIG. **29** shows an exemplary notification-based remedial action method **2900** that may be performed by system **400** and/or any implementation thereof. Method **2900** may be used alone or in combination with any of the other remedial action methods described herein.

(353) At operation **2902**, system **400** determines that a storage system is possibly being targeted by a security threat. This may be performed in any of the ways described herein.

(354) At operation **2904**, system **400** provides a notification in response to the determination that the storage system is possibly being targeted by the security threat. The notification may be in any suitable format. For example, the notification may include a message (e.g., a text message and/or an email), a notification within a user interface used by a user (e.g., an administrator) to manage the storage system, a phone call, and/or any other suitable type of notification as may serve a particular implementation.

(355) FIG. **30** shows an exemplary multi-level remedial action method **3000** that may be performed by system **400** and/or any implementation thereof. Method **3000** may be used alone or in combination with any of the other remedial action methods described herein.

(356) At operation **3002**, system **400** performs a first security threat detection process with respect

to a storage system. The first security threat detection process may include any of the security threat detection processes described herein.

(357) At decision **3004**, system **400** determines, based on the first security threat detection process, whether the storage system is a possible target of a security threat. If system **400** determines that the storage system is not a possible target of security threat based on the first security threat detection process (“No” at decision **3004**), system **400** continues to perform the first security threat detection process at operation **3002**.

(358) However, if system **400** determines, based on the first security threat detection process, that the storage system is a possible target of a security threat (“Yes” at decision **3004**), system **400** may perform a first remedial action (operation **3006**). The first remedial action may include any of the remedial actions described herein.

(359) System **400** may also perform a second security threat detection process with respect to the storage system (operation **3008**). The second security threat detection process may be configured to provide higher confidence threat detection than the first security threat detection process.

Operations **3006** and **3008** may be performed concurrently or sequentially as may serve a particular implementation.

(360) Based on the results of the second security threat detection process, system **400** may either confirm that the storage system is possibly being targeted by the security threat (“Yes” at decision **3010**) or determine that the storage system is not being targeted by the security threat (“No” at decision **3010**), if system **400** determines that the storage system is not being targeted by the security threat (“No” at decision **3010**), system **400** may revert back to performing the first security threat detection process (which may require less resources to perform than the second security threat detection process). However, if system **400** confirms that the storage system is possibly being targeted by the security threat (“Yes” at decision **3010**), system **400** may perform a second remedial action at operation **3012**. The second remedial action may include any of the remedial actions described herein.

(361) In some examples, the second remedial action is different than the first remedial action. For example, the first remedial action may include providing a notification to an administrator of the storage system that the storage system is possibly being targeted by a security threat. If the second security threat detection process confirms this, system **400** may perform a more comprehensive remedial action (the second remedial action), such as creating and/or locking down one or more recovery datasets that may be used to restore corrupted data to an uncorrupted state (such as with the authorization models described herein for ransomware protection snapshots).

(362) Various ways in which the methods and systems of detecting a possible security threat against a storage system and taking one or more remedial actions in response to the security threat are now described.

(363) In some examples, a cloud-based monitoring system implementation of system **400** may provide integrity checks to a storage system or a host that may be used to certify that the storage system or host is running normally and has not been compromised. This may be performed in any suitable manner.

(364) Additionally or alternatively, system **400** may leverage write and deletion protected storage mechanisms to ensure availability of some number of ransomware/corruption protection snapshots, copies, or backups. These are also useful to support legal holds or other related operational purposes.

(365) Additionally or alternatively, system **400** may provide minimum authorization requirements for policy changes (and possibly limits to how already locked down data can be affected by an authorized change in policy) that can be applied to the establishment or configuration of any policies, models, and/or processes described herein. Minimum authorization may require, for example, various combinations of authorization by authenticated operators, administrators, managers, a storage system's vendor, a storage system's selling partner, an AI entity that evaluates

requests, etc. A policy may also stipulate a set of combinations that are allowed to change the policy. Allowed combinations may require, for example, at least a minimum number of managers as well as either multiple entities within the storage system vendor or multiple entities within a storage system's selling partner, as well as certification by at least two of several AI entities evaluating the change. Additionally or alternatively, an additional set of managers (and one or more CxO level authenticated users) may override one or more authorizing entities (e.g., a storage system seller or an AI engine) that would need to authorize a change with fewer managers or without CxCO level authorization from an authenticated CxO level user.

(366) In some examples, duration times for recovery datasets or other time-related models described herein may not generally be based on clocks which are subject to external modification, such as time of day clocks. For example, time interval (or time since power-on) clocks can be used (which are often built into CPUs or other hardware) by system **400**, with interval information being persisted periodically so a reboot or failover within a storage system can leverage prior known intervals to ensure that a minimum absolute time has passed since some time or event associated with a protection snapshot or other aspect of a particular model described herein. This may ensure that an external manipulation of time (such as by hijacking an NTP server on a network) cannot be used to speed up automatic deletion activities. Moreover, if system **400** identifies unusual discrepancies between interval-based time measurement and time-of-day clock times, system **400** may flag this as a potential indicator that the storage system is being targeted by a security threat.

(367) In some examples, system **400** may facilitate replication of data (e.g., rule set data) between administrative authorities. This may provide an additional level of protection against inadvertent or malicious modification of such data.

(368) In some examples, system **400** may direct a first storage system to store replication data in a second storage system with a separate implementation such as through the first storage system storing replicated data as files or objects in a second storage system, or otherwise using the second storage system's regular store operations, rather than through a protocol link between identically implemented storage systems. In this manner, bugs which may be used to attack one of the storage systems may be ineffective at attacking the other storage system.

(369) In some examples, the protection methods and systems described herein may be layered in various ways to increase the robustness of the overall system in ensuring that uncorrupted data is available somewhere. For example, system **400** may provide for storing data into a separately implemented storage system under a separate administrative domain which itself keeps a set of snapshots or includes continuous data protection and which is monitored for corruption by a monitoring service. In this scenario, the primary storage system also includes snapshots or continuous data protection (or both) and is also monitored by the monitoring service.

(370) Some cases of potential corruption or ransomware or other attacks may not be detected by automated software but may be noticed or anticipated by humans. For example, a manager or human resources person may have concerns about a disgruntled employee, or someone in information technology may notice some behaviors that do not make sense. In such cases, a user may provide a user input command for system **400** to direct a storage system to create provisional or locked down ransomware/corruption protection snapshots that may otherwise have been created by policy or by software. In a continuous data protection store, this may result in a set of backward looking locked-down snapshots and recover points, as well. This may also result in a temporary change, with lesser authorization requirements, to increase the rate of creating protection snapshots or to increase the time limits in policies before they can be deleted.

(371) In addition to support for human operators, an API may also be provided by system **400** for creation of ransomware/corruption protection snapshots or for locking down recent snapshots or for managing the creation of provisional protection snapshots (or any other type of provisional ransomware recovery structure) and their conversion to fully protected snapshots. Then, for example, additional security software such as network or server traffic monitors, or software

interacting with software threat analyzer services, may also trigger creation and management of combinations of creation of ransomware/corruption protection snapshots, provisional protection snapshots, and conversions of provisional protection snapshots into full protection snapshots. Such an API may also trigger increases in time limits before protection snapshots can be deleted.

(372) In some examples, a storage system may require certification from a certain number of monitoring services or monitoring service endpoints (such as at least two, or a majority of several such services or endpoints) for the storage system to delete discretionary ransomware protection snapshots and checkpoints or to alter their policy to reduce the period of time they will be retained or to alter the period of time or the amount of activity needed to determine that the provisional detection does or does not rise to the level that the discretionary snapshots will be automatically released or will automatically be converted into full ransomware protection snapshots.

(373) Any of the operations described herein may be performed using a machine learning model, such as any of the machine learning models described herein.

(374) These and other embodiments associated with detecting possible security threats against a storage system, performing remedial actions, managing recovery datasets, and/or otherwise dealing with security threats against a storage system are described more fully in U.S. patent application Ser. No. 16/711,060, filed Dec. 11, 2019, U.S. patent application Ser. No. 16/916,903, filed Jun. 30, 2020 and issued as U.S. Pat. No. 11,341,236, U.S. patent application Ser. No. 16/916,973, filed Jun. 30, 2020, U.S. patent application Ser. No. 16/917,030, filed Jun. 30, 2020, U.S. patent application Ser. No. 16/917,061, filed Jun. 30, 2020, U.S. patent application Ser. No. 17/039,486, filed Sep. 30, 2020, U.S. patent application Ser. No. 17/039,536, filed Sep. 30, 2020, U.S. patent application Ser. No. 17/039,556, filed Sep. 30, 2020, U.S. patent application Ser. No. 17/039,580, filed Sep. 30, 2020 and issued as U.S. Pat. No. 11,500,788, U.S. patent application Ser. No. 17/039,604, filed Sep. 30, 2020, U.S. patent application Ser. No. 17/161,553, filed Jan. 28, 2021 and issued as U.S. Pat. No. 11,520,907, U.S. patent application Ser. No. 17/074,261, filed Oct. 19, 2020, U.S. patent application Ser. No. 17/074,313, filed Oct. 19, 2020, U.S. patent application Ser. No. 17/235,737, filed Apr. 20, 2021, U.S. patent application Ser. No. 17/342,203, filed Jun. 8, 2021, U.S. patent application Ser. No. 17/409,124, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/409,130, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/409,135, filed Aug. 23, 2021, U.S. patent application Ser. No. 17/463,088, filed Aug. 31, 2021, U.S. patent application Ser. No. 17/506,501, filed Oct. 20, 2021, U.S. patent application Ser. No. 17/541,870, filed Dec. 3, 2021, U.S. patent application Ser. No. 17/506,509, filed Oct. 20, 2021, U.S. patent application Ser. No. 17/723,903, filed Apr. 19, 2022, U.S. patent application Ser. No. 17/725,182, filed Apr. 20, 2022, U.S. patent application Ser. No. 17/846,301, filed Jun. 22, 2022, and U.S. patent application Ser. No. 17/980,354, filed Nov. 3, 2022, each of which is incorporated by reference herein.

(375) In some examples, a storage system (e.g., storage system **502**) may be configured to perform a similar block detection process with respect to data that the storage system processes (e.g., data that the storage system receives for storage).

(376) As described herein, blocks of data are similar if they have some number of common substrings or similar patterns. A similarity-based data reduction system implemented by a storage system can identify these blocks that have common substrings or similar patterns in a variety of ways, and then compress those blocks together, or compress them in some other way that shares parts of compression tables, or by describing a subset of the similar blocks as a set of differences from another subset of the similar blocks (for example, based on a single common “template” block). Using such techniques can reduce the amount of physical data stored by a storage system, sometimes substantially.

(377) A special case of similar blocks is identical blocks, which are commonly identified as part of a data deduplication process, e.g., by checksumming, or hashing, whole logical blocks or by identifying a single long representative string within a logical block by using algorithms that identify, for example, a minimum of a rolling hash of some string size over some larger block size

to identify checksums to use for comparison purposes. These checksums may then be compared, such as by looking them up in a table, or performing a sort-merge process, or in some other way, to identify blocks or strings that are identical, usually with some minimum size. This minimum size, or could be a fixed size, and may be in the range of 512 to 8192 bytes, though it could be larger. By focusing on single strings of a minimum size identified through a simple process, such as a single minimum rolling hash or a simple hash of a fixed size, these algorithms can be simpler, in general, than more general similarity-based algorithms. In general, duplicate blocks are reduced simply by multiply referencing a shared copy of the stored identical data, thus avoiding the higher complexity involved in the storing (and eventual garbage collecting) of compression tables or difference maps.

(378) For purposes of the methods and systems described herein, the storage system may simply identify similar blocks, for example because they have at least one identical substring. Methods for identifying similar blocks include the use of MinHash algorithms, “similarity sketches”, and/or similar “feature” identification. It is also sufficient, for purposes of the features described herein, to use algorithms that are probabilistic in nature, as long as the degree of confidence that blocks are similar or identical is high enough that the rate of false positives is sufficiently low that it would not interfere with the ability to identify a probable attack that was altering blocks through encryption or other types of data corruption. For example, a deduplicating storage system may calculate and compare 256-bit secure hashes, which may be the minimum bit size needed to reliably consider two blocks to be actually identical without following up by comparing the data. Some deduplicating storage systems also follow up the hash comparison by comparing the actual data to ensure that the data really is the same, which involves reading the presumed duplicate data so that it can be compared. If only reasonable confidence is needed, following up a hash match by reading and comparing the data may not be needed. Further, for purposes of the methods and systems described herein, the storage system could switch to faster and cheaper 32 bit CRCs and could even store just 16 bits of each CRC (such as the lower 16 bits or the higher 16 bits) to reduce the size of these tables, which may result in a sufficiently low false positive rate.

(379) Accordingly, a similar block detection process may refer to any process used by a storage system to identify blocks of data that are similar to one another and that may be processed to reduce the amount of data stored by a storage system. For example, a similar block detection process may be implemented by a data duplicate detection process used in a data deduplication scheme. Furthermore, the term “similar blocks” may refer to any block of data that is deemed similar or duplicative to another block of data by the similar block detection process.

(380) In some examples, a ransomware attack and/or other malicious action may affect operation of a managed run-time environment configured to access or otherwise be associated with data maintained by a storage system. As used herein, a managed run-time environment may refer to a virtual machine, a containerized application, and/or any other encapsulated and/or isolated virtualized set of compute and/or storage resources operating on a host (e.g., a host server that operates and provides compute resources and provides access to storage resources for a set of such managed run-time environments). In some examples, a managed run-time environment may be defined by a configuration, a backing dataset used for its operating system and application run-time files, and (if running), by currently used CPU and networking resources.

(381) As such, data protection system **400** may be configured to detect an anomaly associated with a managed run-time environment and determine, based on detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat. Based on this determination, data protection system **400** may perform a remedial action associated with the managed run-time environment. For example, data protection system **400** may identify a recovery dataset (e.g., a recovery dataset for data associated with the managed run-time environment and stored by a storage system) generated prior to the anomaly being detected and convert the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity. The remedial action may additionally or alternatively include any of the other

remedial actions described herein.

(382) To illustrate, FIG. 32 shows a method **3200** that may be performed by data protection system **400**. Method **3200** may be performed alone or in combination with any of the other methods described herein.

(383) At operation **3202**, data protection system **400** may detect an anomaly associated with a managed run-time environment. This may be performed in any of the ways described herein.

(384) For example, data protection system **400** may detect the anomaly by determining a value of a metric associated with the managed run-time environment and determining that the value of the metric differs by more than a threshold amount from a baseline value of the metric associated with the managed run-time environment.

(385) The metric may be any suitable metric associated with the managed run-time environment. For example, the metric may include CPU usage by the managed run-time environment.

Additionally or alternatively, the metric may include storage system access operations performed by the managed run-time environment with respect to a storage system that stores data used by and/or otherwise associated with the managed run-time environment.

(386) In some examples, the baseline value to which the metric is compared may be representative of a historical average value for the metric over a predetermined period of time for the managed run-time environment. For example, the historical average value may be an average CPU usage during a particular time period (e.g., during business hours on a weekday, during night hours, during a weekend or holiday, over the course of a lifetime of the managed run-time environment, etc.). As another example, the historical average value may be an average storage system access operation pattern and/or amount during a particular time period. As another example, the historical average value may be an average network usage pattern and/or amount during a particular time period.

(387) Additionally or alternatively, the baseline value may be representative of a value of the metric for an additional managed run-time environment that shares at least one attribute with the managed run-time environment. For example, if the managed run-time environment is a virtual machine, the baseline value may be representative of a value of the metric for another virtual machine that resides on the same host as the virtual machine, executes similar functions as the virtual machine, and/or otherwise has one or more common attributes with the virtual machine.

(388) To illustrate, machine learning and/or other types of processing may identify patterns of CPU usage and/or storage access similar to patterns performed by other managed run-time environments prior to a subsequently detected attack on at least one of those other managed run-time environments, to indicate a relatively early warning of a potential security threat against the managed run-time environment. These patterns may be used to train a machine learning model that can be used by a run-time environment analyzer to identify new security threats that match against the trained model.

(389) In some examples, data protection system **400** may detect the anomaly associated with the managed run-time environment by determining that a host, storage system, and/or other compute environment associated with the managed run-time environment has an anomaly and/or is potentially being targeted by a security threat. For example, if data protection system **400** determines that a storage system that stores data that is accessed by the managed run-time environment is possibly being targeted by a security threat (e.g., in any of the ways described herein), data protection system **400** may determine that an anomaly is associated with the managed run-time environment.

(390) At operation **3204**, data protection system **400** may determine, based on detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat. For example, if run-time CPU usage for the managed run-time environment is greater than a historical average CPU usage for the managed run-time environment, data protection system **400** may determine that the managed run-time environment is possibly being targeted by a security threat. In

some examples, multiple metrics (e.g., CPU usage, network usage, etc.) may be compared to multiple baseline values to generate a higher confidence level determination that the managed run-time environment is possibly being targeted by a security threat.

(391) At operation **3206**, data protection system **400** may perform, based on the determining that the managed run-time environment is possibly being targeted by the security threat, a remedial action associated with the managed run-time environment. The remedial action may include any of the remedial actions described herein.

(392) For example, FIG. **33** shows operations that may be performed by data protection system **400** as part of performing a remedial action with respect to a managed run-time environment. Method **3300** may be performed alone or in combination with any of the other methods described herein.

(393) At operation **3302**, data protection system **400** may identify a recovery dataset (e.g., a snapshot) generated prior to the anomaly being detected at operation **3202**. The recovery dataset may be for data associated with the managed run-time environment and stored by a storage system. For example, the data associated with the managed run-time environment may include data accessed by the managed run-time environment (e.g., data stored in volumes used by the managed run-time environment). Additionally or alternatively, the data associated with the managed run-time environment may include a backing dataset representative of run-time resources that enable operation of the managed run-time environment. For example, the backing dataset may include data representative of one or more executable files that implement the managed run-time environment and/or any other data representative of run-time resources that enable operation of the managed run-time environment. The backing dataset may be distributed, in some examples, across a plurality of hosts and/or storage resources.

(394) In some examples, the recovery dataset generated prior to the anomaly being detected may include a provisional protection recovery dataset, which, as described herein, may be configured to be automatically deleted after a predetermined period of time (e.g., a few days or weeks).

(395) At operation **3304**, data protection system **400** may convert the recovery dataset into a protected recovery dataset. The protected recovery dataset (also referred to as a fully protected snapshot herein) may be configured such that it cannot be deleted without a specific instruction provided by an authorized entity (e.g., an administrator associated with the data protection system **400** and/or the storage system). In this manner, the recovery dataset may be saved and used for a data recovery operation as may serve a particular implementation.

(396) Data protection system **400** may convert the recovery dataset into a protected recovery dataset in any suitable manner. For example, data protection system **400** may update one or more settings and/or metadata associated with the recovery dataset to convert the recovery dataset into the protected recovery dataset.

(397) In some examples, after converting the recovery dataset into a protected recovery dataset, data protection system **400** may continue monitoring the managed run-time environment for anomalies and/or otherwise perform operations to determine whether the managed run-time environment is possibly being targeted by the security threat. If such analysis determines that the managed run-time environment is no longer possibly being targeted by the security threat, data protection system **400** may convert the protected recovery dataset back into a recovery dataset that can be deleted and/or otherwise modified without the specific instruction provided by the authorized entity.

(398) As mentioned, the remedial action performed by data protection system **400** may include any of the remedial actions described herein. For example, the remedial action may further include generating an additional recovery dataset in response to determining that the managed run-time environment is possibly being targeted by the security threat. Additionally or alternatively, the remedial action may further include disabling the managed run-time environment (e.g., until a remedial action may be performed).

(399) In some examples, the data protection system **400** may be implemented by a controller within

the storage system that stores data associated with the managed run-time environment. Additionally or alternatively, data protection system **400** may be implemented by a computing system (e.g., a cloud-based monitoring system) communicatively coupled, by way of a network, to the storage system and/or to a host on which the managed run-time environment resides.

(400) In the preceding description, various exemplary embodiments have been described with reference to the accompanying drawings. It will, however, be evident that various modifications and changes may be made thereto, and additional embodiments may be implemented, without departing from the scope of the invention as set forth in the claims that follow. For example, certain features of one embodiment described herein may be combined with or substituted for features of another embodiment described herein. The description and drawings are accordingly to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A method comprising: detecting, by a data protection system, an anomaly associated with a managed run-time environment; determining, by the data protection system based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; disabling, by the data protection system based on the determining that the managed run-time environment is possibly being targeted by the security threat, the managed run-time environment until a remedial action associated with the managed run-time environment can be performed; performing, by the data protection system based on the determining that the managed run-time environment is possibly being targeted by the security threat, the remedial action, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity; determining, by the data protection system after the performing the remedial action, that the managed run-time environment is no longer possibly being targeted by the security threat; and converting, by the data protection system based on the determining that the managed run-time environment is no longer possibly being targeted by the security threat, the protected recovery dataset back into the recovery dataset.
2. The method of claim 1, wherein the detecting the anomaly associated with the managed run-time environment comprises: determining a value of a metric associated with the managed run-time environment; and determining that the value of the metric differs by more than a threshold amount from a baseline value of the metric associated with the managed run-time environment.
3. The method of claim 2, wherein the metric comprises central processing unit (CPU) usage by the managed run-time environment.
4. The method of claim 2, wherein the metric is representative of storage system access operations performed by the managed run-time environment with respect to the storage system.
5. The method of claim 2, wherein the baseline value is representative of a historical average value for the metric over a predetermined period of time for the managed run-time environment.
6. The method of claim 2, wherein the baseline value is representative of a value of the metric for an additional managed run-time environment that shares at least one attribute with the managed run-time environment.
7. The method of claim 1, wherein the data associated with the managed run-time environment comprises data accessed by the managed run-time environment.
8. The method of claim 1, wherein the data associated with the managed run-time environment comprises a backing dataset representative of run-time resources that enable operation of the managed run-time environment.
9. The method of claim 1, wherein the performing the remedial action further comprises generating

an additional recovery dataset.

10. The method of claim 1, wherein the managed run-time environment comprises at least one of a virtual machine or a containerized application.

11. The method of claim 1, wherein the data protection system is implemented by a controller within the storage system.

12. The method of claim 1, wherein the data protection system is implemented by a computing system communicatively coupled to the storage system by way of a network.

13. A system comprising: a memory storing instructions; and one or more processors communicatively coupled to the memory and configured to execute the instructions to perform a process comprising: detecting an anomaly associated with a managed run-time environment; determining, based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; disabling, based on the determining that the managed run-time environment is possibly being targeted by the security threat, the managed run-time environment until a remedial action associated with the managed run-time environment can be performed; performing, based on the determining that the managed run-time environment is possibly being targeted by the security threat, the remedial action, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity; determining, after the performing the remedial action, that the managed run-time environment is no longer possibly being targeted by the security threat; and converting, based on the determining that the managed run-time environment is no longer possibly being targeted by the security threat, the protected recovery dataset back into the recovery dataset.

14. The system of claim 13, wherein the detecting the anomaly associated with the managed run-time environment comprises: determining a value of a metric associated with the managed run-time environment; and determining that the value of the metric differs by more than a threshold amount from a baseline value of the metric associated with the managed run-time environment.

15. The system of claim 14, wherein the metric comprises central processing unit (CPU) usage by the managed run-time environment.

16. The system of claim 14, wherein the metric is representative of storage system access operations performed by the managed run-time environment with respect to the storage system.

17. The system of claim 14, wherein the baseline value is representative of a historical average value for the metric over a predetermined period of time for the managed run-time environment.

18. The system of claim 14, wherein the baseline value is representative of a value of the metric for an additional managed run-time environment that shares at least one attribute with the managed run-time environment.

19. A non-transitory computer-readable medium storing instructions that, when executed, direct a processor of a computing device to perform a process comprising: detecting an anomaly associated with a managed run-time environment; determining, based on the detecting the anomaly, that the managed run-time environment is possibly being targeted by a security threat; disabling, based on the determining that the managed run-time environment is possibly being targeted by the security threat, the managed run-time environment until a remedial action associated with the managed run-time environment can be performed; performing, based on the determining that the managed run-time environment is possibly being targeted by the security threat, the remedial action, the performing the remedial action comprising: identifying a recovery dataset generated prior to the detecting the anomaly, the recovery dataset for data associated with the managed run-time environment and stored by a storage system, and converting the recovery dataset into a protected recovery dataset that cannot be deleted without a specific instruction provided by an authorized entity; determining, after the performing the remedial action, that the managed run-time

environment is no longer possibly being targeted by the security threat; and converting, based on the determining that the managed run-time environment is no longer possibly being targeted by the security threat, the protected recovery dataset back into the recovery dataset.
