

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12411996
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Kumar; Manoj et al.

Hardware-based implementation of secure hash algorithms

Abstract

A processor includes a register file and an execution unit. The execution unit includes a hash circuit including at least a state register, a state update circuit coupled to the state register, and a control circuit. Based on a hash instruction, the hash circuit receives from the register file and buffers within the state register a current state of a message being hashed. The state update circuit performs state update function on contents of the state register, where performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register. Following completion of all of the plurality of iterative rounds of processing, the execution unit stores contents of the state register to the register file as an updated state of the message.

Inventors: Kumar; Manoj (Yorktown Heights, NY), Mueller; Silvia Melitta (St. Ingbert, DE), Chatterjee; Debapriya (Austin, TX), Fricke; Niels (Herrenberg, DE), Ekanadham; Kattamuri (Mohegan Lake, NY), Boersma; Maarten J. (Holzgerlingen, DE), Berkers; Martijn Diede (Boeblingen, DE)

Applicant: International Business Machines Corporation (Armonk, NY)

Family ID: 87571155

Assignee: International Business Machines Corporation (Armonk, NY)

Appl. No.: 17/884704

Filed: August 10, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20240061961 A1	Feb. 22, 2024

Publication Classification

Int. Cl.: G06F21/72 (20130101); H04L9/06 (20060101)

U.S. Cl.:

CPC G06F21/72 (20130101); H04L9/0643 (20130101); H04L2209/12 (20130101)

Field of Classification Search

USPC: 713/190

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7295671	12/2006	Snell	N/A	N/A
7702100	12/2009	Han et al.	N/A	N/A
7783037	12/2009	Bong	N/A	N/A
8042025	12/2010	Gopal	N/A	N/A
8194854	12/2011	Gueron	N/A	N/A
8340280	12/2011	Gueron	N/A	N/A
8913740	12/2013	Gueron	N/A	N/A
10129018	12/2017	Satpathy	N/A	G06F 21/602
10142098	12/2017	Suresh	N/A	N/A
10256971	12/2018	Gueron	N/A	N/A
10305680	12/2018	Gomes	N/A	N/A
10313108	12/2018	Suresh	N/A	H04L 9/0643
10346343	12/2018	Suresh	N/A	N/A
10348506	12/2018	Greiner	N/A	N/A
10432393	12/2018	Gueron	N/A	N/A
10581593	12/2019	Gomes	N/A	N/A
10877753	12/2019	Bradbury	N/A	N/A
10924282	12/2020	Brostrom	N/A	N/A
11121856	12/2020	Satpathy	N/A	N/A
11128443	12/2020	Gueron	N/A	N/A
2002/0066014	12/2001	Dworkin	N/A	N/A
2003/0135530	12/2002	Parthasarathy	N/A	N/A
2004/0202317	12/2003	Demjanenko	N/A	N/A
2004/0252831	12/2003	Uehara	N/A	N/A
2004/0255130	12/2003	Henry	N/A	N/A
2005/0089160	12/2004	Crispin	N/A	N/A
2008/0240423	12/2007	Gueron et al.	N/A	N/A
2009/0052659	12/2008	Gueron	N/A	N/A
2009/0141887	12/2008	Yap et al.	N/A	N/A
2010/0049986	12/2009	Watanabe	N/A	N/A
2011/0231636	12/2010	Olson	N/A	N/A
2011/0246548	12/2010	Yen	N/A	N/A
2011/0255689	12/2010	Bolotov	N/A	N/A
2015/0043729	12/2014	Gopal	N/A	N/A
2015/0098563	12/2014	Gulley et al.	N/A	N/A

2016/0119126	12/2015	Shay	N/A	N/A
2017/0134163	12/2016	Suresh	N/A	N/A
2017/0373836	12/2016	Rarick	N/A	N/A
2018/0122271	12/2017	Ghosh	N/A	N/A
2019/0197821	12/2018	Hutchinson-Kay et al.	N/A	N/A
2019/0205093	12/2018	Suresh et al.	N/A	N/A
2019/0286443	12/2018	Solomatnikov	N/A	G06F 9/3848
2019/0319782	12/2018	Ghosh	N/A	H04L 9/50
2019/0386815	12/2018	Satpathy	N/A	N/A
2020/0117811	12/2019	Ghosh	N/A	G09C 1/00
2020/0134234	12/2019	Lemay	N/A	N/A
2021/0152330	12/2020	Satpathy	N/A	N/A
2021/0203504	12/2020	Brandt	N/A	N/A
2022/0206958	12/2021	LeMay	N/A	N/A
2023/0269076	12/2022	Brandt	N/A	N/A
2024/0012811	12/2023	Dai	N/A	G06F 16/2255
2024/0015004	12/2023	Chatterjee et al.	N/A	N/A
2024/0053963	12/2023	Mueller et al.	N/A	N/A
2024/0053989	12/2023	Kumar et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1658550	12/2004	CN	N/A
101349968	12/2008	CN	N/A
112152785	12/2019	CN	N/A
113485751	12/2020	CN	N/A
14490631	12/2003	EP	N/A
4569404	12/2024	EP	N/A
4569724	12/2024	EP	N/A
4569725	12/2024	EP	N/A
200536331	12/2004	TW	N/A
201135477	12/2010	TW	N/A
201203108	12/2011	TW	N/A
201332329	12/2012	TW	N/A
201519623	12/2014	TW	N/A
201636829	12/2015	TW	N/A
201717573	12/2016	TW	N/A
202201165	12/2021	TW	N/A

OTHER PUBLICATIONS

Taiwan IPO, P202201322TWo1 Office Action, Jun. 14, 2024, 20 pages (English Translation). cited by applicant

Taiwan IPO, P202201845TW01 Office Action, Jun. 12, 2024, 4 pages (English Translation). cited by applicant

Bertoni, Guido et al., “KangarooTwelve: fast hashing based on Keccak-p.” IACR Cryptology ePrint Archive (2018), 19 pages. cited by applicant

Y. Akiya et al., “SHA-3-LPHP: Hardware Acceleration of SHA-3 for Low-Power High-Performance Systems,” 2021 IEEE International Symposium on Software Reliability Engineering

Workshops (ISSREW), Wuhan, China, 2021, pp. 393-398, doi: 10.1109/ISSREW53611.2021.00107. cited by applicant

International Searching Authority of European Patent Office, International Search Report and Written Opinion, International Application No. PCT/EP2023/071362, Oct. 18, 2023, 13 pages. cited by applicant

Y. Chen et al., "A programmable Galois Field processor for the Internet of Things," 2017 ACM/IEEE 44th Annual International Symposium on Computer Architecture (ISCA), Toronto, ON, Canada, 2017, pp. 55-68, doi: 10.1145/3079856.3080227. cited by applicant

P202201607TW01 Office Action, Jul. 1, 2024, 4 pages, TW Patent Office (English Translation). cited by applicant

Dworkin, M., "SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions NIST FIPS 202," Aug. 2015, 37 pages, National Institute of Standards and Technology, Gaithersburg, MD. cited by applicant

Adams, A. et al., "Cryptography Acceleration in a RISC-V GPGPU," CARRV 2021, Jun. 17, 2021, 7 pages. cited by applicant

Anonymous, "Method of Early Detection and Halting of Ransomware Attacks," IPCOM000268682D, Ip.com, Feb. 15, 2022, 5 pages. cited by applicant

Anonymous, "A Method to Reduce Stick Table Synchronization Storm Among Load Balancer Cluster," IPCOM000268809D, Mar. 1, 2022, 6 pages, Ip.com. cited by applicant

Anonymous, "Method of Flash Acceleration in Replication for Consistent Hash Ring," Mar. 13, 2016, 9 pages, IPCOM000245503D, Ip.com. cited by applicant

Gulley, S. et al., "Intel SHA Extensions: New Instructions Supporting the Secure Hash Algorithm on Intel Architecture Processors," Jul. 2013, 22 pages, Intel Corporation. cited by applicant

Jang, K. et al., "SSL Shader: Cheap SSL Acceleration With Commodity Processors," 8th USENIX Symposium on Networked Systems Design and Implementation, Mar. 2011, 14 pages, USENIX Association, Boston, MA. cited by applicant

NIST FIPS 180-4 "Secure Hash Standard," Aug. 2015, 37 pages, National Institute of Standards and Technology, Gaithersburg, MD. cited by applicant

NIST FIPS 202 "SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions," Aug. 2015, 37 pages, National Institute of Standards and Technology, Gaithersburg, MD. cited by applicant

TW IPO, Notice of Allowance for P202301322TW01, Sep. 26, 2024, 6 pages (English translation). cited by applicant

Taiwan Intellectual Property Bureau, Notice of Allowance in P202201845TW01 (English translation), Aug. 23, 2024, 6 pages. cited by applicant

Advanced Encryption Standard (AES), Nov. 26, 2001, pp. 1-51, FIPS Pub 197, National Institute of Standards and Technology (NIST), US. cited by applicant

Anonymous, "Method and System for Prevention of Anti-Tampering of Media Content," IPCOM000259235D, Jul. 22, 2019, 4 pages, Ip.com. cited by applicant

Anonymous, "Parallelizable Hashing Algorithm," IPCOM000254123D, Jun. 4, 2018, 2 pages, Ip.com. cited by applicant

Bos, J.W. et al., Performance Analysis of the SHA-3 Candidates on Exotic Multi-Core Architectures, CHES 2010 12th International Workshop, Aug. 17-20, 2010, 15 pages, Santa Barbara, CA. cited by applicant

Dworkin, M., "Recommendation for Block Cipher Modes of Operation: Galois/Counter Mode (GCM) and GMAC," NIST Special Publication SP800-38D, Nov. 2007, 39 pages, National institute for Standards and Technology, Gaithersburg, MD. cited by applicant

Ghosh, Moinak, "Optimizing Rolling Hash Computation Using SIMD Vector Registers," IPCOM000226555D, Apr. 16, 2013, 4 pages, Ip.com. cited by applicant

Gueron, S. et al., "Intel Carry-Less multiplication instruction and its usage for computing GCM

Mode,” May 2010, 76 pages, Intel Corporation. cited by applicant

IEEE standard: 1619-2018, “IEEE Standard for Cryptographic Protection of Data on Block-Oriented Storage Devices,” Oct. 23, 2018, 41 pages, IEEE, New York, NY. cited by applicant

International Searching Authority of European Patent Office, International Search Report and Written Opinion, International Application No. PCT/EP2023/068147, Oct. 4, 2023, 14 pages. cited by applicant

Keller, S. et al., “Secure Hash Algorithm 3 Validation System (SHA3VS),” Jan. 29, 2016, 33 pages, National Institute of Standards and Technology, USA. cited by applicant

McGrew, D. et al., “The Galois/Counter Mode of Operation,” Feb. 2004, 43 pages. cited by applicant

Salehani, Y. et al., “NESHHA-256, New 256-BIT Secure Hash Algorithm,” Cryptology ePrint Archive, Paper 2009/033, 2009, 16 pages, <https://eprint.iacr.org/2009/033>. cited by applicant

Appendix P, List of IBM Applications and Patent Treated as Related, 2 pages. cited by applicant

Pittalia, P. P., A comparative study of hash algorithms in cryptography, International Journal of Computer Science and Mobile Computing, 8(6), 147-152 (2019). cited by applicant

Rachh, Rashmi, et al., “Implementation of AES Key Schedule Using Look-Ahead Technique,” Circuits, Systems, and Signal Processing 33.11(2014): 3663-3670. cited by applicant

TW IPO, Office Action for P202201845TW01, Jan. 24, 2024, 9 pages. cited by applicant

International Searching Authority, Notification of Transmittal of the International Search Report and the Written Opinion of the International Searching Authority, Nov. 11, 2023, 12 pages, International Application No. PCT/EP2023/071370. cited by applicant

International Searching Authority, Notification of Transmittal of the International Search Report and the Written Opinion of the International Searching Authority, Nov. 6, 2023, 12 pages, International Application No. PCT/EP2023/071369. cited by applicant

Primary Examiner: Sarker; Sanchit K

Attorney, Agent or Firm: CANTOR COLBURN LLP

Background/Summary

BACKGROUND OF THE INVENTION

- (1) The present invention relates in general to data processing and, in particular, to efficiently performing secure hash algorithms in hardware.
- (2) An important aspect of data security is the protection of data at rest (e.g., while stored in a data storage device) or data in transit (e.g., during transmission) through encryption. Generally speaking, encryption involves the conversion of unencrypted data (referred to as plaintext) to encrypted data (referred to as ciphertext) through the combination of the plaintext and one or more encryption keys utilizing an encryption function. To recover the plaintext from the ciphertext, the ciphertext is processed by a decryption function utilizing one or more decryption keys. Thus, encryption provides data security by the requirement that an additional secret (i.e., the decryption key(s)) be known by a party prior to that party being able to access the protected plaintext.
- (3) In many implementations, data encryption is performed utilizing software executing on a general purpose processor. While the implementation of encryption in software provides the advantages of being able to select differing encryption algorithms and easily adapting a selected encryption algorithm to use various data lengths, performing encryption in software has the concomitant disadvantage of relatively poor performance. As the volumes of data sets continue to dramatically increase in the era of “big data,” the performance achieved by software-implemented

encryption can be unacceptable when encrypting large messages and/or data sets. The concern with encryption performance also arises due to the increased need to run enterprise applications utilizing encrypted data in order to mitigate the consequence of “hacking” and other cyber-attacks and to ensure regulatory compliance. Accordingly, it is often desirable to provide support for encryption in hardware to achieve improved performance.

BRIEF SUMMARY

(4) The present disclosure appreciates that one class of encryption algorithm for which it would be desirable to provide hardware support is hash functions, including but not limited to the hash functions belonging to the Secure Hash Algorithm (SHA) family of standards. As is known in the art, the SHA family of standards define hash algorithms approved by the National Institute for Standards (NIST) for generating a condensed representation of a message (i.e., a message digest). The SHA family of standards is specified in two Federal Information Processing Standards (FIPS): FIPS 180-4 “Secure Hash Standard” and FIPS 202 “SHA-3 Standard: Permutation-Based Hash and Extendable-Output Functions,” which are incorporated herein by reference. FIPS 180-4 specifies seven hash algorithms, namely, SHA-1 (Secure Hash Algorithm-1) and the SHA-2 family of hash algorithms, including SHA-224, SHA-256, SHA-384, SHA-512, SHA-512/224, and SHA-512/256. FIPS 202 additionally specifies four SHA-3 hash algorithms having a fixed-length output (i.e., SHA3-224, SHA3-256, SHA3-384, and SHA3-512) and two closely related “extendable-output” functions (XOFs) named SHAKE128 and SHAKE256 (where SHAKE is an abbreviation for Secure Hash Algorithm and Keccak). Additional uses for the SHA family of standards (e.g., as a stream cipher, an authenticated encryption system, or a tree hashing scheme) have not yet been adopted as NIST standards.

(5) Given the wide diversity of hash functions and data sizes of hash functions (even within the SHA family of standards), broad support in hardware for hash functions can lead to large areas within a processor floorplan being consumed by circuitry implementing hash functions. As a result, some hardware solutions choose to implement such circuitry separately from a processor core, for example, in a bus-attached Application-Specific Integrated Circuits (ASICs) or accelerator. While offering the possibility of better performance than some software solutions, use of these adjunct circuits remains subject to bus and memory access latencies and messaging overhead, again limiting performance as compared to that achievable within a high-performance processor core. This performance penalty is particularly acute for relatively small messages (e.g., messages fitting within a single message block), which are the majority of SHA messages handled in enterprise servers. The present disclosure addresses these and other design considerations by efficiently implementing hash functions in a processor.

(6) In one embodiment, a processor includes a register file and an execution unit. The execution unit includes a hash circuit including at least a state register, a state update circuit coupled to the state register, and a control circuit. Based on a hash instruction, the hash circuit receives from the register file and buffers within the state register a current state of a message being hashed. The state update circuit performs state update function on contents of the state register, where performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register. Following completion of all of the plurality of iterative rounds of processing, the execution unit stores contents of the state register to the register file as an updated state of the message.

(7) Such a processor can be incorporated into a data processing system including multiple processors, a shared memory, and a system interconnect communicatively coupling the shared memory and the multiple processors. Such a processor can also be tangibly embodied in a design structure in a machine-readable storage device for designing, manufacturing, or testing an integrated circuit.

(8) In one embodiment, a method of data processing in a processor includes fetching, by an

instruction fetch unit, instructions to be executed by the processor. The instructions include a hash instruction. Based on receiving the hash instruction, an execution unit of the processor executes the hash instruction. Executing the hash instruction includes receiving from a register file and buffering, within a state register of the execution unit, a current state of a message being hashed. Execution of the hash instruction also includes performing, in a state update circuit, a state update function on contents of the state register, where performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register. Following completion of all of the plurality of iterative rounds of processing, contents of the state register are stored into the register file as an updated state of the message.

(9) In at least some embodiments, the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function and the state update circuit performs twenty-four rounds of processing each utilizes as an input a respective one of twenty-four round indices.

(10) In one embodiment, the execution unit executes the hash instruction in a squeezing phase of a Secure Hash Algorithm and Keccak (SHAKE) hash algorithm.

(11) In at least some embodiments, the state update function comprises a Secure Hash Algorithm 2 (SHA2) block hash function. In at least some embodiments, the hash circuit further includes an adder configured to add contents of the state register to the current state and return a resulting sum to the register file. In some embodiments, the execution unit further includes a message block register for buffering a message block of the message and a message schedule round circuit coupled to the message block register. The message schedule round circuit performs a plurality of iterative rounds of processing on contents of the message block register and returns a result of each of the plurality of iterative rounds of processing to the message block register. In some embodiments, the state update circuit includes a data path for data words having a first data width and the execution unit is configured, based on the hash instruction indicating a second data width that is narrower than the first data width, to expand data words of a message block of the message to the first data width prior to processing the data words of the message in the state update circuit.

Description

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

- (1) FIG. 1 is a high-level block diagram of a data processing system including a processor in accordance with one embodiment;
- (2) FIG. 2 is a high-level block diagram of a processor core in accordance with one embodiment;
- (3) FIG. 3 is a high-level block diagram of an exemplary execution unit of a processor core in accordance with one embodiment;
- (4) FIG. 4 is a more detailed block diagram of an accelerator unit within a processor core in accordance with one embodiment;
- (5) FIG. 5 is a time-space diagram of message hashing in accordance with the SHA-3 standard;
- (6) FIG. 6 is a time-space diagram of the absorption phase depicted in FIG. 5;
- (7) FIG. 7A is a time-space diagram of the SHA3 permute function illustrated in FIG. 6;
- (8) FIG. 7B is a time-space diagram of the SHA3 round function depicted in FIG. 7A;
- (9) FIG. 8 is a time-space diagram of the SHA3/SHAKE squeezing phase illustrated in FIG. 5;
- (10) FIGS. 9-10 illustrate exemplary formats for a SHA3 hash instruction and bitwise exclusive OR (XOR) instruction, respectively, in accordance with one embodiment;
- (11) FIG. 11 is a high-level block diagram of an exemplary SHA3/SHAKE hash circuit in accordance with one embodiment;
- (12) FIG. 12 is a high-level logical flowchart of an exemplary process by which a processor executes a SHA3 hash instruction in accordance with one embodiment;

- (13) FIG. 13 depicts a time-space diagram of message hashing in accordance with the SHA-2 standard;
- (14) FIG. 14 is a time-space diagram of the SHA2 block hash function illustrated in FIG. 13;
- (15) FIG. 15 illustrates message expansion for a SHA2 hash function having 32-bit words in accordance with an exemplary embodiment;
- (16) FIG. 16 depicts an exemplary format for a SHA2 hash instruction in accordance with one embodiment;
- (17) FIG. 17 is a high-level block diagram of an exemplary SHA2 hash circuit in accordance with one embodiment;
- (18) FIG. 18 is a high-level block diagram of an exemplary update working state circuit from FIG. 17 in accordance with one embodiment;
- (19) FIG. 19 is a high-level block diagram of an exemplary embodiment of a SHA2 sigma circuit as shown in FIG. 18;
- (20) FIG. 20 is a high-level logical flowchart of an exemplary process by which a processor executes a SHA2 hash instruction in accordance with one embodiment;
- (21) FIG. 21A depicts an exemplary unpadded message;
- (22) FIG. 21B illustrates an exemplary padded message;
- (23) FIGS. 22A-22B depict the assembly of chunks of a message block in a narrower first register file and the transfer of the message block to a wider second register file;
- (24) FIGS. 23A-23D illustrate various padding scenarios for SHA3/SHAKE messages;
- (25) FIGS. 24A-24D depict various padding scenarios for SHA2 messages;
- (26) FIG. 25 illustrates an exemplary padding instruction in accordance with one embodiment;
- (27) FIG. 26 depicts an exemplary padding circuit in accordance with one embodiment;
- (28) FIG. 27 illustrates an exemplary circuit for combining end-of-block (EOB) and end-of-message (EOM) bytes with a message in accordance with one embodiment;
- (29) FIG. 28 is a high-level logical flowchart of an exemplary process for padding a message block in accordance with one embodiment; and
- (30) FIG. 29 depicts an exemplary design process in accordance with one embodiment.

DETAILED DESCRIPTION

(31) With reference now to the figures and in particular with reference to FIG. 1, there is illustrated a high-level block diagram of a data processing system **100** in accordance with one embodiment. In some implementations, data processing system **100** can be, for example, a server computer system (such as one of the POWER series of servers available from International Business Machines Corporation), a mainframe computer system, a mobile computing device (such as a smartphone or tablet), a laptop or desktop personal computer system, or an embedded processor system.

(32) As shown, data processing system **100** includes one or more processors **102** that process instructions and data. Each processor **102** may be realized as a respective integrated circuit having a semiconductor substrate in which integrated circuitry is formed, as is known in the art. In at least some embodiments, processors **102** can generally implement any one of a number of commercially available processor architectures, for example, POWER, ARM, Intel x86, NVidia, Apple silicon, etc. In the depicted example, each processor **102** includes one or more processor cores **104** and cache memory **106** providing low latency access to instructions and operands likely to be read and/or written by processor cores **104**. Processors **102** are coupled for communication by a system interconnect **110**, which in various implementations may include one or more buses, switches, bridges, and/or hybrid interconnects.

(33) Data processing system **100** may additionally include a number of other components coupled to system interconnect **110**. These components can include, for example, a memory controller **112** that controls access by processors **102** and other components of data processing system **100** to a system memory **114**. In addition, data processing system **100** may include an input/output (I/O) adapter **116** for coupling one or I/O devices to system interconnect **110**, a non-volatile storage

system **118**, and a network adapter **120** for coupling data processing system **100** to a communication network (e.g., a wired or wireless local area network and/or the Internet).

(34) Those skilled in the art will additionally appreciate that data processing system **100** shown in FIG. **1** can include many additional non-illustrated components. Because such additional components are not necessary for an understanding of the described embodiments, they are not illustrated in FIG. **1** or discussed further herein. It should also be understood, however, that the enhancements described herein are applicable to data processing systems and processors of diverse architectures and are in no way limited to the generalized data processing system architecture illustrated in FIG. **1**.

(35) Referring now to FIG. **2**, there is depicted a high-level block diagram of an exemplary processor core **200** in accordance with one embodiment. Processor core **200** may be utilized to implement any of processor cores **104** of FIG. **1**.

(36) In the depicted example, processor core **200** includes an instruction fetch unit **202** for fetching instructions within one or more streams of instructions from storage **230** (which may include, for example, cache memories **106** and/or system memory **114** from FIG. **1**). In a typical implementation, each instruction has format defined by the instruction set architecture of processor core **200** and includes at least an operation code (opcode) field specifying an operation (e.g., fixed-point or floating-point arithmetic operation, vector operation, matrix operation, logical operation, branch operation, memory access operation, cryptographic operation, etc.) to be performed by processor core **200**. Certain instructions may additionally include one or more operand fields directly specifying operands or implicitly or explicitly referencing one or more registers storing source operand(s) to be utilized in the execution of the instruction and one or more registers for storing destination operand(s) generated by execution of the instruction. Instruction decode unit **204**, which in some embodiments may be merged with instruction fetch unit **202**, decodes the instructions retrieved from storage **230** by instruction fetch unit **202** and forwards branch instructions that control the flow of execution to branch processing unit **206**. In some embodiments, the processing of branch instructions performed by branch processing unit **206** may include speculating the outcome of conditional branch instructions. The results of branch processing (both speculative and non-speculative) by branch processing unit **206** may, in turn, be utilized to redirect one or more streams of instruction fetching by instruction fetch unit **202**.

(37) Instruction decode unit **204** forwards instructions that are not branch instructions (often referred to as “sequential instructions”) to mapper circuit **210**. Mapper circuit **210** is responsible for the assignment of physical registers within the register files of processor core **200** to instructions as needed to support instruction execution. Mapper circuit **210** preferably implements register renaming. Thus, for at least some classes of instructions, mapper circuit **210** establishes transient mappings between a set of logical (or architected) registers referenced by the instructions and a larger set of physical registers within the register files of processor core **200**. As a result, processor core **200** can avoid unnecessary serialization of instructions that are not data dependent, as might otherwise occur due to the reuse of the limited set of architected registers by instructions proximate in program order.

(38) Still referring to FIG. **2**, processor core **200** additionally includes a dispatch circuit **216** configured to ensure that any data dependencies between instructions are observed and to dispatch sequential instructions as they become ready for execution. Instructions dispatched by dispatch circuit **216** are temporarily buffered in an issue queue **218** until the execution units of processor core **200** have resources available to execute the dispatched instructions. As the appropriate execution resources become available, issue queue **218** issues instructions from issue queue **218** to the execution units of processor core **200** opportunistically and possibly out-of-order with respect to the original program order of the instructions.

(39) In the depicted example, processor core **200** includes several different types of execution units for executing respective different classes of instructions. In this example, the execution units

include one or more fixed-point units **220** for executing instructions that access fixed-point operands, one or more floating-point units **222** for executing instructions that access floating-point operands, one or more load-store units **224** for loading data from and storing data to storage **230**, and one or more vector-scalar units **226** for executing instructions that access vector and/or scalar operands. In a typical embodiment, each execution unit is implemented as a multi-stage pipeline in which multiple instructions can be simultaneously processed at different stages of execution. Each execution unit preferably includes or is coupled to access at least one register file including a plurality of physical registers for temporarily buffering operands accessed in or generated by instruction execution.

(40) Those skilled in the art will appreciate that processor core **200** may include additional unillustrated components, such as logic configured to manage the completion and retirement of instructions for which execution by execution units **220-226** is finished. Because these additional components are not necessary for an understanding of the described embodiments, they are not illustrated in FIG. 2 or discussed further herein.

(41) With reference now to FIG. 3, there is illustrated a high-level block diagram of an exemplary execution unit of a processor **102** in accordance with one embodiment. In this example, a vector-scalar unit **226** of processor core **200** is shown in greater detail. In the embodiment of FIG. 3, vector-scalar unit **226** is configured to execute multiple different classes of instructions that operate on and generate different types of operands. For example, vector-scalar unit **226** is configured to execute a first class of instructions that operate on vector and scalar source operands and that generate vector and scalar destination operands. Vector-scalar unit **226** executes instructions in this first class of instructions in functional units **302-312**, which in the depicted embodiment include an arithmetic logic unit/rotation unit **302** for performing addition, subtraction, and rotation operations, a multiply unit **304** for performing binary multiplication, a divide unit **306** for performing binary division, a cryptography unit **308** for performing cryptographic functions, a permute unit **310** for performing operand permutations, and binary-coded decimal (BCD) unit **312** for performing decimal mathematical operations. The vector and scalar source operands on which these operations are performed and the vector and scalar destination operands generated by these operations are buffered in the physical registers of an architected register file **300**.

(42) In this example, vector-scalar unit **226** is additionally configured to execute a second class of instructions that cause hashing functions to be performed. Vector-scalar unit **226** executes instructions in this second class of instructions in an accelerator unit **314**. The operands on which these hash functions are performed and the operands generated by these hash functions are buffered and accumulated in a wide vector register file **316**, which may include, for example, 1024-bit wide physical registers.

(43) In operation, instructions are received by vector-scalar unit **226** from issue queue **218**. If an instruction is in the first class of instructions (e.g., vector-scalar instructions), the relevant source operand(s) for the instruction are accessed in architected register file **300** utilizing the mapping between logical and physical registers established by mapper circuit **210** and then forwarded with the instruction to the relevant one of functional units **302-312** for execution. The destination operand(s) generated by that execution are then stored back to the physical register(s) of architected register file **300** determined by the mapping established by mapper circuit **210**. If, on the other hand, the instruction is in the second class of instructions (e.g., hash instructions), the instruction is forwarded to accelerator unit **314** for execution with respect to operand(s) buffered in specified registers of wide vector register file **316**.

(44) Referring now to FIG. 4, there is depicted a more detailed block diagram of accelerator unit **314** of FIG. 3 in accordance with one embodiment. Accelerator unit **314** includes circuitry for performing, in hardware, a variety of hash functions, including, for example, one or more hash functions defined by the SHA family of standards. In the depicted example, the hash circuitry of accelerator unit **314** includes at least SHA3/SHAKE hash circuit **400** as described below in greater

detail with reference to FIG. 11 and a SHA2 hash circuit 402 as described below in greater detail with reference to FIG. 17. Accelerator unit 314 additionally includes a single instruction multiple data (SIMD) exclusive OR (XOR) circuit 404 employed in performing a SHA3/SHAKE hash of a message, as discussed further below. Finally, accelerator unit 314 includes a data transfer circuit 406 that transfers data (e.g., messages to be hashed and message digests) between the memory system (e.g., cache memory 106 and system memory 114) and wide vector register file 316.

(45) With reference now to FIG. 5, there is a time-space diagram of a process 500 of message hashing in accordance with the SHA-3 standard. As is known in the art, the SHA-3 standard (i.e., FIPS 202) employs a sponge construction based on a wide random function or random permutation. According to this sponge construction, a message 502 of any arbitrary length (possibly many megabytes) is first processed in an input phase, referred to as the SHA3 absorption phase 504 in sponge terminology. SHA3 absorption phase 504, which is described in greater detail below with reference to FIG. 6, is identical for both SHA3 and SHAKE hash functions. SHA3 absorption phase 504 produces a 1600-bit last absorption state 610, which is then processed in an output phase (referred to as the SHA3/SHAKE squeezing phase 506 in sponge terminology) to generate a message digest 508. SHA3/SHAKE squeezing phase 506, which is described in detail below with reference to FIG. 8, operates differently for SHA3 hash functions and SHAKE hash functions. In particular, SHA3/SHAKE squeezing phase 506 generates fixed-length message digests 508 for the various SHA3 hash functions, but generates variable-length message digests 508 for SHAKE hash functions.

(46) Table I, below, summarizes attributes of four SHA3 hash functions and two SHAKE hash functions defined by the SHA-3 standard and listed in the first column. In Table I, the second column summarizes the size in bits (r) of the message blocks into which the SHA3 absorption phase 504 subdivides variable-length message 502. Message block size r is an integer multiple of bytes in length, and the first message block of each message is byte-aligned. The third column of Table I summarizes the size in bits (d) of the message digest 508 output by SHA3/SHAKE squeezing phase 506. It should again be noted that SHAKE-128 and SHAKE-256, unlike the SHA3 hash functions, generate variable length digests of length d'. As noted in the fourth column of Table I, for each hash function specified by the SHA-3 standard, the last absorption state 610 is 1600 bits in length. The fifth column of Table I specifies the differing values of c, a number of lower-order bits passed between iterations of a SHA3 state permute function during the SHA3/SHAKE squeezing phase 506 (see, e.g., FIG. 8). Finally, the sixth column of Table I specifies that each iteration of the SHA3 state permute function employs 24 rounds of permutation per message block (see, e.g., FIG. 7A). In updates to the SHA-3 standard or in non-standard implementations, the number of rounds of permutation may be varied, for example, by decreasing the required number of permutations (e.g., to 12).

(47) TABLE-US-00001

Function	r (bits)	d (bits)	c (bits)	Iterations	State (bits)	Permutation (rounds)
SHA3-224	112	224	1600	448	24	24
SHA3-256	1088	256	1600	512	24	24
SHA3-384	832	384	1600	768	24	24
SHA3-512	576	512	1600	1024	24	24
SHAKE-128	1344	d'	1600	256	24	24
SHAKE-256	1088	d'	1600	512	24	24

(48) Referring now to FIG. 6, there is depicted a time-space diagram of the SHA3 absorption phase 504 depicted in FIG. 5. As shown, SHA3 absorption phase 504 receives, as input, a message 502 of any arbitrary length. As shown at block 600, message 502 is padded to obtain a length that is an integer multiple of r bits. In many prior art implementations, this padding is accomplished via a high-latency, computationally expensive memory-to-memory move of the entire message 502. In some other prior art implementations, the SHA hash software routine pads message blocks after loading the message block into SIMD registers using a sequence of conventional SIMD instructions. Although these prior art techniques can be utilized herein to perform padding, as described in detail below with reference to FIGS. 21A to 27, this padding can instead be efficiently performed by hardware in processor registers (e.g., wide vector register file 316) in accordance

with the disclosed inventions through execution of a padding instruction. Padding message **502** through execution of a padding instruction also allows the padding to be applied to the end of message **502** in a temporally overlapping manner with the processing of message blocks in SHA3 absorption phase **504**.

(49) In SHA3 absorption phase **504**, each of n (n a positive integer) message blocks of length r composing the padded message is extracted and then zero-extended in the trailing low-order bits to form n 1600-bit extended message blocks **602**. The first message block, that is, message block 1 **602**, forms the input of a SHA3 state permute function **604** defined by the SHA-3 standard. As described below with reference to FIGS. **9** and **11**, in accordance with one aspect of the disclosed inventions, SHA3 state permute function **604** is performed in hardware through the execution of a SHA3 hash instruction. The 1600-bit state output of SHA3 state permute function **604** forms a first input of a 1600-bit bitwise XOR function **606**, which takes as a second input a next 1600-bit extended message block **602** of the padded message. The result of the bitwise XOR function **606** forms the input of a next iteration of the SHA3 state permute function **604**. As shown, this process continues iteratively for each of the message blocks **602** until a final iteration of SHA3 state permute function **604** generates and outputs a 1600-bit last absorption state **610**, as previously noted in the description of FIG. **5**.

(50) With reference now to FIG. **7A**, there is illustrated a time-space diagram of the SHA3 permute function **604** illustrated in FIG. **6**. SHA3 permute function **604** takes in a 1600-bit input and then processes the 1600-bit input in conjunction with a SHA-3 standard-specified round index 0 **702** in a first of 24 rounds of a SHA3 round function **704**. This process continues iteratively, with each subsequent round of processing in a SHA3 round function **704** receiving as inputs the 1600-bit output of the preceding SHA3 round function **704** and the relevant SHA3 standard-specified round index **702** (which is a constant). After the 24 rounds of processing within SHA3 state permute function **604** are complete, SHA3 state permute function **604** outputs a 1600-bit state, which either serves as an input to a bitwise XOR function **606**, or in the case of the final iteration of SHA3 state permute function **604** within SA3 absorption phase **504**, constitutes the last absorption state **610** that serves as the input of SHA3/SHAKE squeezing phase **506**.

(51) Referring now to FIG. **7B**, there is depicted a time-space diagram of the SHA3 round function **704** depicted in FIG. **7A**. As shown, SHA3 round function **704** includes a SHA-3 standard-specified sequence of functions, including in order, five functions referred to in the SHA-3 standard by the Greek letters θ (theta), ρ (rho), π (pi), χ (chi), and $\{ \text{umlaut over } (\cdot \text{Math.}) \}$ (iota). The theta function receives and processes the 1600-bit input to round function **704**, and the output of each other function except the iota function feeds the next sequential function. Finally, the iota function processes the output of the chi function and the relevant round index **702** to produce the 1600-bit output of the given iteration of SHA3 round function **704**. In the prior art, performing round function **704** can take up to 80 cycles utilizing two single instruction multiple data (SIND) vector pipelines. In accordance with one aspect of the inventions disclosed herein, round function **704** can be completed in a single cycle of a processor core **104** utilizing the SHA3/SHAKE hash circuit **400** of FIG. **11** described below.

(52) With reference now to FIG. **8**, there is illustrated a time-space diagram of the SHA3/SHAKE squeezing phase **506** illustrated in FIG. **5**. As previously described, SHA3/SHAKE squeezing phase **506** receives as input the 1600-bit last absorption phase **610** produced by SHA3 absorption phase **504**. To produce a message digest **508** for any of the SHA3 functions defined by the SHA-3 standard, SHA3/SHAKE squeezing phase **506** first extracts the first r high-order bits of last absorption state **610** to form result block 1 **800**. A truncate function **802** then truncates the r bits of result block 1 **800** to retain the high-order d bits forming message digest **508**.

(53) To produce a message digest for one of the SHAKE functions defined by the SHA-3 standard, the r bits of result block 1 **800** form the r high-order bits of the input of a truncate function **804**. These r high-order bits are concatenated with $n-1$ additional r -bit result blocks **800**, each of which

is formed from the r high-order bits of the output of an iteration of SHA3 state permute function **604** as previously described with respect to FIG. 7A. Each SHA3 state permute function **604** of SHA3/SHAKE squeezing phase **506** receives a 1600-bit input (i.e., $r+c=1600$) and generates a 1600-bit output, which, except for the last iteration of SHA3 state permute function **604**, feeds a subsequent iteration of SHA3 state permute function **604**. Truncate function **804** truncates the $r \times n$ input bits to obtain a message digest **508** having a user-specified length of d' bits.

(54) Referring now to FIGS. **9-10**, there are illustrated exemplary formats for a SHA3 hash instruction **900** and bitwise exclusive OR (XOR) instruction **1000**, respectively, in accordance with one embodiment. In an exemplary embodiment, accelerator unit **314** is configured, responsive to receipt of SHA3 hash instruction **900**, to perform a SHA3/SHAKE state permute function in hardware utilizing SHA3/SHAKE hash circuit **400** and, responsive to receipt of bitwise XOR instruction **1000** to perform a 1024-bit bitwise XOR of specified operands utilizing SIND XOR circuit **404**.

(55) In the illustrated embodiment, SHA3 hash instruction **900** includes an opcode field **902** that specifies the particular architecture-specific operation code for a SHA3/SHAKE permute function. SHA3 hash instruction **900** additionally includes one or more register fields **904**, **906** for specifying the registers within wide vector register file **316** for the source and destination operands of the SHA3/SHAKE state permute function. For example, in one implementation, SHA3 hash instruction **900** includes a single register field **904** that specifies the first of a pair of adjacent 1024-bit registers that buffer the 1600-bit source operand and, following completion of the SHA3/SHAKE permute function, the 1600-bit destination operand (which overwrites the source operand). In an alternative implementation, SHA3 hash instruction **900** includes two register fields **904**, **906** for specifying separate pairs of 1024-bit source and destination registers (in which case the destination operand does not overwrite the source operand).

(56) As noted above, in future updates to the SHA-3 standard or in non-standard implementations, it can be desirable to control the number of rounds of permutation applied by the SHA3 state permute function **604**. In such embodiments, the number of rounds SHA3 hash instruction **900** can include a field that either sets the number of rounds of permutation directly or references a register specifying the number of rounds of permutation.

(57) FIG. **10** depicts an exemplary embodiment in which bitwise XOR instruction includes an opcode field **1002** that specifies the particular architecture-specific operation code for a 1024-bit bitwise XOR function. Bitwise XOR instruction **1000** additionally includes three register fields **1004**, **1006**, and **1008** for separately specifying 1024-bit registers within wide vector register file **316** for buffering two 1024-bit source operands and one 1024-bit destination operand.

(58) Now that the SHA3 and SHAKE hash functions and exemplary instructions for implementing portions of these hash functions have been explained, pseudocode for performing an exemplary SHA3 hash function in hardware is presented. In the pseudocode that follows, the following registers are referenced:

(59) TABLE-US-00002 $R_r \leftarrow$ Block length in bytes $R_L \leftarrow$ Message length in bytes // It is assumed $R_L \geq R_r$ & the first block is not padded $R_a \leftarrow$ Starting address of the message $R_b \leftarrow$ Address of message digest resulting from the hash function $R_d \leftarrow$ Message digest length in bytes $X_s \leftarrow$ SHA3 state // Wide vector register pair $X_m \leftarrow$ Message block // Wide vector register pair
Given these registers, pseudocode for any of the SHA3 (not SHAKE) hash function can be expressed as follows:

(60) TABLE-US-00003 $X_s = \text{loadlength}(R_a, R_r)$ // Load first message block of message and initialize state $X_s = \text{sha3hash}(X_s)$ // Execute SHA3 hash instruction to perform permute on // first message block $R_L = R_r$ // Decrement length of unprocessed portion of message $R_a += R_r$ // Increment pointer to next message block in message While ($R_L \geq R_r$) // Enter loop for processing each remaining message block, // except for last message block of message { $X_m = \text{loadlength}(R_a, R_r)$ // Load next message block $X_s = \text{wide_xor}(X_s, X_m)$ // Execute bitwise XOR

instruction to combine state and // current message block Xs = sha3hash(Xs) // Execute SHA3 hash instruction to perform permute on // current message block RL = RL - Rr // Decrement length of unprocessed portion of message Ra += Rr // Increment pointer to next message block } Xm = loadlength(Ra, RL) // Load last message block, if present (RL can be zero) Xm = sha3_padding(Xm, RL, sha3-type) // Execute padding instruction to pad // message based on remaining message // length and SHA3 function Xs = wide_xor(Xs, Xm) // Execute bitwise XOR instruction to combine state and // last message block Xs = sha3 hash(Xs) // Execute SHA3 hash instruction to perform permute on // last message block and produce last absorption state Store_length(Xs, Rb, Rd) // In SHA3 squeezing phase, truncate the last absorption // state to form the message digest by storing the leading Rd // bytes of Xs to memory at address Rb

(61) With reference now to FIG. 11, there is illustrated a high-level block diagram of an exemplary SHA3/SHAKE hash circuit **400** suitable for executing a SHA3 hash instruction **900** in accordance with one embodiment. As shown, SHA3/SHAKE hash circuit **400** includes two 1024-bit two-input multiplexers **1100a**, **1100b**, two 1024-bit state registers **1102a**, **1102b**, a SHA3 round circuit **1106**, and a control circuit **1110** that controls operation of SHA3/SHAKE hash circuit **400** in response to a SHA3 hash instruction **900**.

(62) Input multiplexer **1100a** has a first input coupled to receive the high-order 1024 bits of a 1600-bit input state from a first register of a register pair in wide vector register file **316** identified by SHA3 hash instruction **900** and a second input coupled to receive the high-order 1024 bits of the 1600-bit round feedback from SHA3 round circuit **1106**. Input multiplexer **1100b** is similarly structured, with a first input coupled to receive, from a second register in the instruction-specified register pair in wide vector register file **316**, a 1024-bit value including the low-order 576 bits of a 1600-bit input state and a second input coupled to SHA3 round circuit **1106** to receive a 1024-bit value including the low-order 576 bits of the 1600-bit round feedback. Control logic **1110** within SHA3/SHAKE hash circuit **400** provides unillustrated select signals to input multiplexers **1100a**, **1100b** to cause input multiplexers **1100a**, **1100b** to select the values present at their first inputs prior to SHA3 round 0 and to select the values present at their second inputs following each of SHA3 round() through SHA3 round 23. The values output by input multiplexers **1100a**, **1100b**, which are buffered in state registers **1102a**, **1102b**, respectively, together form the 1600-bit round input value of SHA3 round circuit **1106**, which is configured to perform a SHA3 round function **704** on the round input value as previously described with reference to FIGS. 7A-7B.

(63) Control circuit **1110** is further configured to sequence SHA3 round circuit **1106** through each of the 24 rounds required by the SHA-3 standard utilizing the correct round indices specified by the SHA-3 standard. Following conclusion of round 23, state registers **1102a**, **1102b** will respectively hold the high-order 1024 bits and low-order 576 bits of the 1600-bit output state. Control circuit **1110** is further configured, once the output state is obtained, to assert unillustrated select signals to cause an output multiplexer **1108** to write the high-order bits and low-order bits of the 1600-bit output state from state registers **1102a**, **1102b**, respectively, to an instruction-specified register pair in wide vector register file **316** in two consecutive cycles (assuming wide vector register file **316** has a single write port).

(64) Referring now to FIG. 12, there is depicted a high-level logical flowchart of an exemplary process for executing a SHA3 hash instruction **900** in accordance with one embodiment. For ease of understanding, the process of FIG. 12 is described with reference to exemplary SHA3/SHAKE hash circuit **400** of FIG. 11.

(65) The process of FIG. 12 begins at block **1200** and then proceeds to block **1202**, which illustrates SHA3/SHAKE hash circuit **400** receiving a SHA3 hash instruction **900** specifying a operand register pair within wide vector register file **316**. In response to receipt of the SHA3 hash instruction **900**, control circuit **1110** causes the contents of the operand register pair to be read out from wide vector register file **316** and loaded into state registers **1102a**, **1102b** via input multiplexers **1100a**, **1100b** (block **1204**). Control circuit **1110** additionally initializes an internal

round counter to 0 (block **1206**).

(66) The process then proceeds from block **1206** to block **1208**, which illustrates control circuit **1110** directing SHA3 round circuit **1106** to perform an iteration of the SHA3 round function **704** utilizing the round input buffered in state registers **1102a**, **1102b** and the appropriate SHA-3 standard-specified round index. Control circuit **1110** additionally increments the round counter (block **1208**). The result of the processing of SHA3 round circuit **1106** is returned by input multiplexers **1100a**, **1100b** to state registers **1102**, **1102b**. As indicated at block **1210**, control logic **1110** causes SHA3 round circuit **1106** to perform 24 rounds of processing specified by the SHA-3 standard utilizing the appropriate round indices. When the 24 rounds of processing are complete, control circuit **1110** asserts appropriate select signals to cause output multiplexer **1108** to store the 1600-bit state buffered in state registers **1102a**, **1102b** (as zero-extended in the low-order bits to form two 1024-bit values) into the operand register pair within wide vector register file **316** specified by SHA3 hash instruction **900** (block **1214**). Thereafter, the process of FIG. 12 ends at block **1216**.

(67) With reference now to FIG. 13, there is illustrated a time-space diagram of message hashing in accordance with the SHA-2 standard (FIPS 180-4), which in the embodiment of FIG. 4 is performed by SHA2 hash circuit **402**. Table II, below, summarizes attributes of six SHA2 hash functions defined by the SHA-2 standard and listed in the first column. In Table II, the second column summarizes the message block size (r) in bits. Message block size r is an integer multiple of bytes in length, and the first message block of the message is byte-aligned. The third column of Table II summarizes the fixed size in bits (d) of the message digest resulting from each SHA2 hash function. The fourth column of Table II species the size in bits of the state for each SHA2 hash function, and the fifth column of Table II indicates the number of rounds of processing (i.e., 64 or 80) employed in each SHA2 hash function (see, e.g., FIG. 14). Finally, the sixth column of Table II specifies the word size, in bits, employed for each SHA2 hash function. Note that for all variants, the state size is 8 times the word size (i.e., comprises 8 words), and the size of the message block is 16 times the size of a word (i.e., comprises 16 words). As described below, in accordance with one aspect of the disclosed inventions, SHA2 hash functions employing a 32-bit word size and SHA2 hash functions employing a 64-bit word size are processed along the same dataflow by virtue of the message expansion applied to words of SHA2-224 and SHA2-256 hash functions, as described below with reference to FIG. 15.

(68) TABLE-US-00004

TABLE II	Message block size r (bits)	Message digest size d (bits)	State size (bits)	Word size w (bits)	Rounds
SHA2-224	112	224	1120	32	64
SHA2-256	128	256	1280	32	64
SHA2-384	384	384	3840	128	80
SHA2-512	1024	512	4096	128	80
SHA2-224/256	112	224	1120	64	80
SHA2-256/224	128	256	1280	64	80

(69) As shown in FIG. 13, a SHA2 hash function **1300** receives, as one input, a message **1302** of any arbitrary length (e.g., possibly megabytes in length). As shown at block **1304**, message **1302** is padded to obtain a length that is an integer multiple of r bits. As discussed above with reference to FIG. 6, this padding can be efficiently performed by hardware in processor registers (e.g., wide vector register file **316**) rather than through a memory move through execution of a padding instruction. Padding message **1302**, and specifically, padding the last message block(s) of the message **1302**, through execution of a padding instruction also allows the padding to be applied to the end of message **1302** in a temporally overlapping manner with processing of message blocks by SHA2 hash function **1300**. Each of n (n a positive integer) message blocks of length r (where $r=16 \times w$) composing the padded message produced by block **1304** is extracted to form one of n $16 \times w$ -bit message blocks **1306**.

(70) In addition to message **1302**, SHA2 hash function **1300** also receives as an input a SHA-2-specified constant value of $8 \times w$ bits. As is known in the art, this constant value, which can be accessed from architected register file **300**, varies between SHA2 hash function sand forms an $8 \times w$ -bit initial state **1308**. Initial state **1308** and the first message block, that is, message block 1

1306, form two inputs of a SHA2 block hash function **1310** defined by the SHA-2 standard. As described below with reference to FIGS. **16** and **17**, in accordance with one aspect of the disclosed inventions, SHA2 block hash function **1310** is performed in hardware through the execution of a SHA2 hash instruction. The $8 \times w$ -bit state output by SHA2 block hash **1** function **1310** forms a first input of a SHA2 block hash function **2** **1310**, which takes as a second input a next $16 \times w$ -bit message block **2** **1306**. The result of the SHA2 block hash function **2** **130** forms the input of a next iteration of the SHA2 block hash function **1310**. As shown, this process continues iteratively for each of the message blocks **602** until a final nth iteration of SHA2 block hash function **1310** generates and outputs a $8 \times w$ -bit last state, which is truncated by truncate function **1312** to produce a message digest **1314** of d bits.

(71) Referring now to FIG. **14**, there is depicted a time-space diagram of the SHA2 block hash function **1310** illustrated in FIG. **13**. SHA2 block hash function **1310** takes in a $16 \times w$ -bit message block **1306** and, as shown at block **1420**, initializes a $16 \times w$ -bit message schedule for the message block **1306**. SHA2 block hash function **1310** then processes the $16 \times w$ -bit message schedule through n rounds of processing in message schedule round function **1400**, with the $16 \times w$ -bit output of each of rounds 1 to n-2 serving as the input to a next round of message schedule processing.

(72) As shown, SHA2 block hash function **1310** also receives as an input the $8 \times w$ -bit current hash state (i.e., either initial state **1308** or the output of a previous SHA2 block hash function **1310**). As indicated at block **1406**, SHA2 block hash function **1310** splits this $8 \times w$ -bit current hash state into 8 w-bit variables a to h. SHA2 block hash function **1310** then processes the current hash state through n rounds of processing by update round function **1404**. Initial update round 0 **1404** takes as additional inputs a SHA-2 specified w-bit round key 0 **1402** and the w high order bits of a $16 \times w$ -bit initialization **1420** of the message schedule. Each succeeding iteration of update round function **1404** takes as inputs the state generated by the preceding iteration of update round function **1404**, the w high-order bits of the $16 \times w$ -bit output of the corresponding iteration of message schedule round function **1400**, and a SHA-2 specified w-bit round key **1402**. The hash state output by update round function n-1 **1404** is added to the input hash state by a $8 \times w$ -bit carry-propagating add function **1410** to generate a next hash state.

(73) With reference now to FIG. **15**, message expansion for a SHA2 hash function in accordance with an exemplary embodiment is illustrated. As noted above with reference to Table II and FIG. **13**, embodiments of the disclosed inventions preferably support processing of SHA2 hash functions of differing word sizes w along common data paths by expanding the message words and initial hash states of those SHA2 hash functions employing smaller word sizes. This expansion can be performed, for example, at blocks **1304** and **1308** of FIG. **13**. FIG. **15** illustrates a specific example in which each of the sixteen 32-bit words **1502** of a SHA2-224 or SHA2-256 input message **1500** is expanded to form a corresponding one of sixteen 64-bit double words **1506** of an output message **1504**. In this example, each 64-bit double word **1506** is formed by concatenating a 32-bit word of input message **1500** in the high-order half of the 64-bit double word **1506** with a 32-bit zero word **1508** in the low-order half of the double word **1506**. The resulting output message **1504** can then be processed by a SHA2 hash circuit in the same manner as messages employing 64-bit words.

(74) Referring now to FIG. **16**, there is depicted an exemplary format for a SHA2 hash instruction **1600** in accordance with one embodiment. In an exemplary embodiment, accelerator unit **314** is configured, responsive to receipt of SHA2 hash instruction **1600**, to perform a SHA2 block hash function **1310** in hardware utilizing SHA2 hash circuit **402**.

(75) In the illustrated embodiment, SHA2 hash instruction **1600** includes an opcode field **1602** that specifies the particular architecture-specific operation code for a SHA2 block hash function. SHA2 hash instruction **1600** additionally includes one or more operand register fields **1604**, **1606** for specifying the operand registers within wide vector register file **316** for the source and destination operands of the SHA2 block hash function. For example, in one implementation, SHA2 hash instruction **1600** includes a register field **1604** that specifies the 1024-bit register that buffers the

input current hash state and, following completion of the SHA2 block hash function, buffers the output current hash state (which overwrites the input current hash state). In addition, SHA2 hash instruction **1600** includes register field **1606** that buffers a current message block to be processed. SHA2 hash instruction **1600** further includes a mode field **1608** indicating whether the SHA2 hash function to be performed employs 32-bit or 64-bit words.

(76) Now that the SHA2 hash functions and an exemplary instruction for implementing portions of the SHA2 hash functions have been explained, pseudocode for performing an exemplary SHA2 hash function (i.e., SHA2-512) in hardware is presented. In the SHA2-512 hash function, each message block is 1024 bits in length, and the hash state and message digest are each 512 bits in length. In the pseudocode that follows, the following registers are referenced:

(77) TABLE-US-00005 $RL \leftarrow$ message length in bits $RL \leftarrow$ message length in bytes; assumed to be ≥ 128 bytes so first message block is not padded $Ra \leftarrow$ starting address of message $Ri \leftarrow$ address of initial state $Rb \leftarrow$ address of the message digest resulting from hash function $Rd \leftarrow$ message digest length in bytes $Xs \leftarrow$ SHA2 state // wide vector register $Xm \leftarrow$ current message block // wide vector register

(78) Given these registers, pseudocode for performing the SHA2-512 hash function can be expressed as follows:

(79) TABLE-US-00006 $Xs = \text{load}(Ri, 64)$ // Load initial state of 64 bytes $Xm = \text{load}(Ra, 128)$ // Load first (full) message block $Xs = \text{sha2hash}(Xs, Xm, 64\text{-bit})$ // Execute SHA2 hash instruction to perform block // hash function $RL - = 128$ // Decrement message length to be processed $Ra += 128$ // Advance pointer to next message block While ($RL > = 128$) // Loop through remaining message blocks, except // last message block { $Xm = \text{load}(Ra, 128)$ // Load next message block (full size) $Xs = \text{sha2hash}(Xs, Xm, 64\text{-bit})$ // Execute SHA2 hash instruction to perform block // hash function $RL - = 128$ // Decrement message length to be processed $Ra += 128$ // Advance pointer to next message block } $Xm = \text{loadlength}(Ra, RL)$ // Load last message block, if present (RL can be zero) $Xm = \text{sha2_EOM_pad}(Xm, RL)$ // Append SHA2 EOM byte to end of message // block If ($RL > 111$) then // If padding crosses two message blocks, { $Xs = \text{sha2hash}(Xs, Xm, 64\text{-bit})$ // execute SHA2 hash instruction to perform block $Xm = \text{force-to-zero}$ // hash function and zero last message block } $Xm = \text{sha2_EOB_pad}(Xm, Ri)$ // Insert EOB in last block of the padded message $Xs = \text{sha2hash}(Xs, Xm, 64\text{-bit})$ // Execute SHA2 hash instruction to perform block // hash function on last message block Store($Xs, Rb, 64$) // Truncate state to leading 64 bytes of Xs to obtain // message digest and store to memory at address // Rb

(80) With reference now to FIG. 17, there is illustrated a high-level block diagram of an exemplary embodiment of SHA2 hash circuit **402** of FIG. 4 that is suitable for executing a SHA2 hash instruction **1600**. As shown, SHA2 hash circuit **402** includes a 512-bit two-input state multiplexer **1702a**, a 1024-bit two-input message multiplexer **1702b**, a 512-bit state register **1704a**, a 1024-bit message block register **1704b**, an update working state circuit **1708**, a message schedule round circuit **1710**, and a control circuit **1720** that controls operation of SHA2 hash circuit **402** in response to a SHA2 hash instruction **1600**.

(81) In this example, the first input of state multiplexer **1702a** is coupled to receive, from the register in wide vector register file **316** specified by register field **1604** of SHA2 hash instruction **1600**, a current hash state held in the 512 high-order bits of the register. The second input of state multiplexer **1702a** is coupled to the output of update working state circuit **1708**. Message multiplexer **1702b** is similarly configured, with a first input coupled to receive a message block from a register in wide vector register file **316** specified by register field **1606** of SHA2 hash instruction **1600** and a second input coupled to receive the 1024-bit round feedback from message schedule round circuit **1710**. Control logic **1720** within SHA2 hash circuit **400** provides unillustrated select signals to multiplexers **1702a**, **1702b** to cause multiplexers **1702a**, **1702b** to select the values present at the first inputs prior to update round 0 function **1404** and to select the values present at the second inputs following each of update round 0 function through SHA2 block

hash n function. The values output by multiplexer **1702a**, **1702b** are temporarily buffered in state and message block registers **1704a**, **1704b**, respectively. The message block buffered in message block register **1704b** forms the input of message schedule round circuit **1710**, which implements message schedule round function **1400** of FIG. **14**. The 64 high-order bits from the message block register **1704b** and the 512-bit state in state register **1704a** form the two inputs of update working state circuit **1708**, which is configured to perform an update round function **1404** as previously described with reference to FIG. **14**.

(82) Control circuit **1720** is further configured to sequence update working state circuit **1708** through each of the n rounds utilizing the correct round index specified by the SHA-2 standard. Following conclusion of the last round $n-1$, state register **1704a** will hold a 512-bit hash state. Control circuit **1720** is further configured, once the output hash state is obtained, to cause single instruction multiple data (SIMD) adder **1712** to add the hash state from state register **1704a** with the input hash state read from wide vector register file **316** and to store the result, which is the next hash state, back to wide vector register file **316**, as described above with respect to add function **1410** of FIG. **14**. Those skilled in the art will appreciate that in differing implementations, SIMD adder **1712** can either be implemented as a dedicated component of SHA2 hash circuit **402** or as a separate pipeline that can be shared, for example, by multiple hash circuits.

(83) Referring now to FIG. **18**, there is depicted a more detailed block diagram of an exemplary update working state circuit **1708** from FIG. **17** in accordance with one embodiment. In this embodiment, the 512-bit state buffered within state register **1704a**, which is received as one input of update working state circuit **1708**, is split into eight 64-bit variables referred to in the SHA-2 standard as variables a to h , as shown at block **1800**. Update working state circuit **1708** includes two sigma function circuits, namely, SHA2 sigma0 circuit **1802** and SHA2 sigma1 circuit **1806**, as well as a SHA2 MA circuit **1804**, and a SHA2 CH circuit **1808**, which each performs a respective function defined by the SHA-2 standard. Update working state circuit **1708** additionally includes three 64-bit adders **1810**, **1812**, and **1814**. SHA2 sigma0 circuit **1802** applies a sigma function with $n(n_1, n_2, n_3)=(28, 34, 39)$ and $m(m_1, m_2, m_3)=(2, 13, 22)$ to variable a to produce a first input of adder **1812**. Variables a , b , and c are processed by SHA2 MA circuit **1804** to produce a second input of adder **1812**. SHA2 sigma1 circuit **1806** applies a sigma function with $n(n_1, n_2, n_3)=(14, 18, 41)$ and $m(m_1, m_2, m_3)=(6, 13, 22)$ to variable e to produce a first input among the five inputs of adder **1810**. Variables e , f and g are processed by SHA2 CH circuit **1808** to produce a second input of adder **1810**. Adder **1810** adds to these two inputs the relevant round key, the round message block, and variable d to produce a sum that forms a first input of adder **1814** and a third input of adder **1812**.

(84) Update working state circuit **1708** generates a 512-bit result state **1816** composed of eight 64-bit variables a' to h' . Variable a' of result state **1816** is formed by the output of adder **1812**, variables b' , c' , and d' are respectively formed of variables a , b , and c of input state **1800**, and variables f' , g' , and h' are respectively formed of variables e , f and g of input state **1800**. Remaining variable e' is formed by the sum of the output of adder **1810** and variable d of input state **1800**.

(85) It should be noted that the 32-bit to 64-bit expansion of the words of SHA-2 messages described above with reference to FIG. **15** does not impact the design of (is transparent to) SHA2 MA circuit **1804**, SHA2 CH circuit **1808**, and modular adders **1812**, **1814**. The trailing zero expansion of SHA2 messages employing 32-bit words only impacts SHA2 sigma circuits **1802**, **1806**, as described in greater detail below with reference to FIG. **19**.

(86) FIG. **19** is a more detailed block diagram of an exemplary embodiment of a SHA2 sigma circuit **1900**, which can be utilized to implement SHA2 sigma0 circuit **1802** and SHA2 sigma1 circuit **1806** of FIG. **18**. SHA2 sigma circuit **1900** receives a 64-bit input variable **1902** including 32 high-order bits (bits 0 to 31) and 32 low-order bits (bits 32 to 63).

(87) SHA2 sigma circuit **1900** includes a 64-bit rotate circuit **1904a** that rotates the 64-bit input variable **1902** by n_1 bits (i.e., 28 bits for SHA2 sigma0 circuit **1802** and 14 bits for SHA2 sigma1

circuit **1806**) to obtain a first 64-bit input of multiplexer **1910a**. SHA2 sigma circuit **1900** additionally includes a 32-bit rotate circuit **1906a** that rotates the 32 high-order bits of input variable **1902** by m1 bits (i.e., 2 bits for SHA2 sigma0 circuit **1802** and 6 bits for SHA2 sigma 1 circuit **1806**) to obtain, when concatenated with the 32 low-order bits of input variable **1902**, a second 64-bit input of multiplexer **1910a**. Multiplexer **1910a** selects between its first input and second input based on a mode signal determined by mode field **1608** of the relevant SHA2 hash instruction **1600**. That is, multiplexer **1910a** selects the first input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 64-bit words and selects the second input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 32-bit words.

(88) SHA2 sigma circuit **1900** additionally includes a 64-bit rotate circuit **1904b** that rotates the 64-bit input variable **1902** by n2 bits (i.e., 34 bits for SHA2 sigma0 circuit **1802** and 18 bits for SHA2 sigma1 circuit **1806**) to obtain a first 64-bit input of multiplexer **1910b**. SHA2 sigma circuit **1900** also includes a 32-bit rotate circuit **1906b** that rotates the 32 high-order bits of input variable **1902** by m2 bits (i.e., 13 bits for both SHA2 sigma0 circuit **1802** and SHA2 sigma 1 circuit **1806**) to obtain, when concatenated with the 32 low-order bits of input variable **1902**, a second 64-bit input of multiplexer **1910b**. Multiplexer **1910b** selects between its first input and second input based on the mode signal. In particular, multiplexer **1910b** selects the first input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 64-bit words and selects the second input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 32-bit words.

(89) SHA2 sigma circuit **1900** also includes a 64-bit rotate/shift circuit **1908a** that rotates and shifts the 64-bit input variable by n3 bits (i.e., 39 bits for SHA2 sigma0 circuit **1802** and 41 bits for SHA2 sigma1 circuit **1806**) to obtain a first 64-bit input of multiplexer **1910c**. SHA2 sigma circuit **1900** additionally includes a 32-bit rotate/shift circuit **1908b** that rotates and shifts the 32 high-order bits of input variable **1902** by m3 bits (i.e., 22 bits for both SHA2 sigma0 circuit **1802** and SHA2 sigma 1 circuit **1806**) to obtain, when concatenated with the 32 low-order bits of input variable **1902**, a second 64-bit input of multiplexer **1910c**. Multiplexer **1910c** selects between its first input and second input based on the mode signal. As with multiplexers **1910a**, **1910b**, multiplexer **1910c** selects the first input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 64-bit words and selects the second input if the mode signal indicates mode field **1608** is set to indicate a SHA2 hash function that utilizes 32-bit words.

(90) The 64-bit outputs of multiplexers **1910a**, **1910b**, and **1910c** form the inputs of a three-input 64-bit bitwise XOR circuit **1912**, which performs a bitwise XOR on its three inputs to generate a 64-bit output **1914**. Those skilled in the art will appreciate that in some embodiments of SHA2 sigma circuit **1900**, the functions of rotate circuits **1904a-1904b** and **1906a-1906b** and rotate/shift circuits **1908a-1908b** can be implemented by appropriate wiring, allowing the SHA2 sigma circuit **1900** to be realized with three multiplexers **1910a-1910c** and a 3-way bitwise XOR circuit **1912** and without explicit rotate and shift circuitry.

(91) Referring now to FIG. 20, there is depicted a high-level logical flowchart of an exemplary process for executing a SHA2 hash instruction **1600** in accordance with one embodiment. For ease of understanding, the process of FIG. 20 is described with reference to the exemplary embodiment of SHA2 hash circuit **402** illustrated in FIGS. 17-19.

(92) The process of FIG. 20 begins at block **2000** and then proceeds to block **2002**, which illustrates SHA2 hash circuit **402** receiving a SHA2 hash instruction **1600** specifying a particular SHA2 mode (i.e., 32-bit or 64-bit word size), as well as state and message block registers within wide vector register file **316**. In response to receipt of the SHA2 hash instruction **1600**, control circuit **1720** causes the 512-bit state and 1024-bit message block to be read out from wide vector register file **316** and loaded into state register **1704a** and message block register **1704b**, respectively, via multiplexers **1702a**, **1702b** (block **2002**). Control circuit **1720** additionally

initializes an internal round counter to 0 (block **2004**).

(93) The process then proceeds from block **2004** to block **2006**, which illustrates control circuit **1720** directing message schedule round circuit **1710** to perform an iteration of the message schedule round function **1400** utilizing the message block buffered in message block register **1704b**. In addition, control circuit **1720** directs update working state circuit **1708** to perform an iteration of update round function **1404** based on the appropriate round index, the 64 high-order bits of message block register **1704b**, and the input hash state from state register **1704a**. The results of the processing of update working state circuit **1708** and message schedule round circuit **1710** are returned by multiplexers **1702a**, **1702b** to registers **1704a**, **1704b**, respectively. Control circuit **1110** additionally advances the round counter. At block **2010**, control logic **1720** determines by reference to the round counter whether SHA2 hash circuit **402** has performed the last rounds of processing specified by the SHA-2 standard. As noted in Table II, SHA2 hash circuit **402** performs 64 rounds of processing for SHA2 hash functions employing 32-bit words and 80 rounds of processing for SHA2 hash functions employing 64-bit words. If control circuit **1720** determines at block **2010** that at least one additional round of processing remains to be performed, the process returns to block **2006**, which has been described. However, in response to determining at block **2010** that all rounds of processing are complete, control circuit **1720** causes the prior state to again be read from wide vector register file **316** and to be added by SIMD adder **1712** to the final state buffered in state register **1704a** (block **2012**). Control circuit **1720** then directs storage of the resulting next state back into wide vector register file **316** (block **2014**). Thereafter, the process of FIG. **20** ends at block **2016**.

(94) As discussed above with reference to block **600** of FIG. **6** and block **1304** of FIG. **13**, messages processed by the SHA2 and SHA3 hashing functions are padded to create messages having lengths that are even multiples of block length r bits. FIG. **21A** depicts an exemplary unpadded message **2100**, which has a total length of L bits and includes n message blocks. Of these, the first $n-1$ message blocks include r bits, but the final message block n includes k bits, where $k < r$. As shown in FIG. **21B**, in the general case, message **2100** is padded by appending $r-k$ padding bits to the end of message block n , resulting in n message blocks all of a length of r bits.

(95) The content of padding bits appended to obtain the padded message can vary depending on the hash function under consideration. For example, in the SHA2 and SHA3/SHAKE hash algorithms discussed herein the padding bits will include bytes marking both the end of the unpadded portion of the message (i.e., an end-of-message (EOM) marker) and the end of last block of the padded message (i.e., an end-of-block (EOB) marker). As explained further below, in some cases, padding bits including the EOM and EOB markers can all be included within the message block containing the final message byte; in other cases, addition of the padding bits can require appending an additional message block to the message. In either case, the disclosed inventions preferably perform message padding in processor registers through execution of one or more instructions rather than through a high-latency memory move operation that transfers the message between two locations in memory.

(96) In at least some architectures, load-store units **224**, memory controller **112**, and/or system interconnect **110** are not constructed to support data transfers of lengthy data objects (e.g., complete r -bit SHA3/SHAKE and SHA2 message blocks) between system memory **114** and wide vector register file **116**. In such architectures, message blocks are transferred in multiple smaller chunks into a narrower register file and then transferred from the narrower register file into one or more wide vector registers of wide vector register file **316**. For example, FIG. **22A** illustrates an example in which a SHA3/SHAKE message block n is assembled in an architected register file **300** including 256-bit registers r_0 to r_S **301**. In this example, a load-length instruction is executed, for example, by a load-store unit **224** of FIG. **2**, to load five 256-bit chunks of a 1152-bit SHA3-224 message block n into registers r_0 to r_7 and to zero any register bytes not containing message data. Given the 1152-bit length of message blocks in SHA3-224, the message bytes within message

block n will, at most, completely fill registers r0 to r3 plus the leading 128 bits of register r4 (of course, the final message block of an unpadded message may contain fewer than r bits). At least the remaining 128 bits of register r4 and all of registers r5 to r7 can be zeroed, either automatically by execution of the load-length instruction or by execution of standard load instructions. (Filling registers r6 and r7 with zeros is only required for generic SHA3 functions applicable to any of the supported message block lengths.) An additional data transfer instruction may then be executed by data transfer circuit **406** or transfer unit **320** to transfer the contents of registers r0 to r7 into registers R0 and R1 317 of wide vector register file **316**, which includes registers R0 to RT each having an exemplary width of 1024 bits as discussed above. In an alternative implementation, the same result can be achieved by loading four registers **301** within architected register file **300** to buffer chunks n1 to n4 and then reusing the same registers **301** on a subsequent cycle to buffer chunks n5 to n8.

(97) FIG. **22B** depicts a similar example showing the transfer of a 1024-bit SHA2 message block n to a wide vector register **317** in wide vector register file **316** following assembly of the message block in registers **301** of architected register file **300**. In this example, a load-length instruction is executed, for example, by load-store unit **224** of FIG. **2**, to load four 256-bit chunks of a SHA2 message block n into registers r2 to r5 of architected register file **300** and to zero any register bytes not containing message data. An additional data transfer instruction may then be executed by data transfer circuit **406** to transfer the contents of registers r2 to r5 into register R0 **317** of wide vector register file **316**. In an alternative implementation, the same result can be achieved by loading two registers **301** within architected register file **300** to buffer chunks n1 and n2 and then reusing the same registers **301** on a subsequent cycle to buffer chunks n3 to n4.

(98) In at least some preferred embodiments, the process for loading message blocks into wide vector register file **316** given in FIGS. **22A-22B** is performed for all message blocks of a SHA3/SHAKE or SHA2 message, including message block n, which is the last message block of the unpadded message. As explained below, the end of the message can then be padded, at least in part, within wide vector register file **316** through execution of one or more instructions.

(99) FIGS. **23A-23D** depict various padding cases for various lengths of SHA3/SHAKE messages. In accordance with the SHA-3 standard, each message must include EOM padding marking the EOM. Under the SHA3-standard, the EOM padding has a fixed value of x06 for SHA3 hash functions and a fixed value of x1F for SHAKE hash functions. The location of the EOM padding within the padded message varies depending on the message length, which is often unknown at compile time. The SHA-3 standard further mandates that the last byte of each padded message is a fixed-value EOB padding byte.

(100) As shown in FIG. **23A**, if the last message block **2300** of a SHA3/SHAKE message includes more than two bytes that do not contain message data, an EOM padding byte **2302** is inserted into the zeroed byte of the relevant wide vector register **317** immediately following the last message byte **2306**, and an EOB padding byte **2304** is inserted into the zeroed byte of the wide vector register **317** as the last byte of the padded message block.

(101) FIG. **23B** illustrates a similar second case in which last message block **2300'** of the SHA3/SHAKE message includes exactly two zeroed bytes that do not contain message data. In this case, the last two zeroed bytes of last message block **2300'** are replaced with the EOM padding byte **2302** followed by EOB padding byte **2304**.

(102) FIG. **23C** depicts a third case in which the last message block **2300''** of a SHA3/SHAKE message includes only a single zeroed message byte following last message byte **2306**. In this case, execution of a padding instruction as described below causes the EOM and EOB padding values to be ORed together and inserted in the final byte of padded message block **2300''** as EOM/EOB padding byte **2308**.

(103) FIG. **23D** illustrates the final case in which the last message byte **2306** of the SHA3/SHAKE message is the last byte of a message block **2310**. Because message block **2310** does not include

capacity for the required EOM and EOB padding in this case, an additional zeroed message block **2312** is appended to the message (e.g., through execution of a load-length instruction). Into this zeroed message block **2312**, an EOM padding byte **2302** is inserted as the first byte and an EOB padding byte **2304** is inserted as the last byte. It should be noted that, in each of the four cases depicted in FIGS. **23A-23D**, both EOM padding and EOB padding can be advantageously applied by a single padding instruction since both the EOM padding and EOB padding always fall within the same message block. It should also be appreciated that although FIGS. **23A-23D** depict the application of padding to messages including an integer number of message bytes, padding can similarly be applied to bit messages not including an integer number of bytes.

(104) In one embodiment, padding of a SHA3/SHAKE message of arbitrary length as shown in FIGS. **23A-23D** can be implemented utilizing three instructions. These instructions include: (1) a load-length instruction that stages the final message block of the padded message in specified registers **301** in architected register file **300**, (2) a transfer instruction that transfers the message block from the registers **301** in architected register file **300** to one or more wide vector registers **317** in wide vector register file **316** as shown in FIG. **22A**, and (3) a padding instruction that inserts EOM and EOB padding at the appropriate byte location(s) in the final message block of the padded SHA3/SHAKE message held in the wide vector register(s) **317**. Of course, it is possible in alternative implementations to insert the EOM padding and EOB padding into the final message block utilizing two different instructions. However, for single block messages such as those commonly utilized in post-quantum cryptography schemes, the addition of an extra padding instruction increases latency and undesirably reduces hashing performance.

(105) FIGS. **24A-24D** depict various padding cases for various lengths of SHA2 messages. In accordance with the SHA-2 standard, each message must include one EOM padding byte of value **x80** in the byte immediately following the last message byte. The location of the EOM padding byte within the padded message thus varies depending on message length. The SHA-2 standard further mandates that the last two words (i.e., either two 32-bit or two 64-bit words depending on the SHA2 hash function in question (see Table II)) contain EOB padding specifying the length of the unpadded message in bits.

(106) In the first case illustrated in FIG. **24A**, the last message block **2400** of a SHA2 message includes more than two words plus one byte that do not contain message data. In this case, last message block **2400** is padded by insertion of an EOM padding byte **2302** into the zeroed byte of the relevant wide vector register **317** immediately following last message byte **2406** and by insertion of two EOB padding words **2404** as the last two words of last message block **2400**.

(107) FIG. **24B** illustrates a similar second case in which the last message block **2400'** of a SHA2 message includes exactly two words plus one byte that do not contain message data. In this case, last message block **2400'** is padded by insertion of an EOM padding byte **2402** into the zeroed byte of the relevant wide vector register **317** immediately following last message byte **2406** and insertion of two EOB padding words **2404** as the last two words of last message block **2400**.

(108) FIG. **24C** depicts a third case in which the last message block **2400''** of an unpadded SHA2 message includes too few bytes not containing message data to accommodate both a EOM padding byte **2402** and two EOB padding words **2404**. In this case, the SHA2 message is padded by insertion of an EOM padding byte **2402** into the zeroed byte of the relevant wide vector register **317** immediately following last message byte **2406**. Because the EOB padding words **2404** do not fit within message block **2400''**, an additional zeroed message block **2408** is appended to the message (e.g., through execution of a load-length instruction). The EOB padding words **2404** are then inserted as the last two words of message block **2408**.

(109) FIG. **24D** illustrates a fourth case in which the last message byte **2406** of the SHA2 message forms the last byte of a complete message block **2410**. Because message block **2410** does not include capacity for the EOM or EOB padding, an additional zeroed message block **2412** is appended to the SHA2 message. Additional message block **2412** includes EOM padding byte **2302**

as the first byte of message block **2412**, followed by a number of zeroed bytes, and finally two EOB padding words **2404** at the end of message block **2412**.

(110) In one embodiment, padding of a SHA2 message of arbitrary length can be implemented utilizing as few as four instructions. These instructions include: (1) a load-length instruction that places a final message block of the SHA2 message in specified registers **301** in architected register file **300** and that zeroes any register bytes not containing message bytes (2) an insert-word instruction that places two EOB padding words **2404** in the appropriate bytes of a register **301** in architected register file **300** to mark the end of the padded message, (3) a transfer instruction that transfers contents of the registers **301** buffering the message block from architected register file **300** to a wide vector register **317** in wide vector register file **316**, and (4) a padding instruction that inserts EOM padding byte **2402** at the appropriate location in the wide vector register **317**. In this embodiment, execution of the padding instruction inserts EOM padding byte **2402** but not EOB padding words **2404** because (1) the EOM padding byte **2402** and EOB padding words **2404** may be located in different message blocks and (2) EOB padding words **2404** can be efficiently positioned in the appropriate register **301** within architected register file **300** utilizing an existing insert-word instruction. Of course, in an alternative embodiment, both the EOM padding byte **2402** and EOB padding words **2404** can be applied to a SHA2 message block in the registers **301** of architected register file **300**.

(111) With reference now to FIG. **25**, there is illustrated an exemplary padding instruction **2500** in accordance with one embodiment. In at least one embodiment, exemplary padding instruction **2500** may be executed by accelerator unit **314** within data transfer circuit **406** to perform padding for both SHA3/SHAKE and SHA2 message blocks.

(112) In the illustrated example, padding instruction **2500** includes an opcode field **2502** that specifies the architecture-specific operation code for a message padding instruction. Padding instruction additionally includes two register fields **2504**, **2506** for specifying the storage locations of source and destination operands of the padding operation. For example, register1 field **2504** can identify the target wide vector register **317** within wide vector register file **316** that buffers the message block to be padded, and register2 field **2506** can specify the register **301** in architected register file **300** holding the remaining message length in bytes.

(113) Padding instruction **2500** further includes a mode field **2508** providing information utilized to pad the message. In one exemplary embodiment, mode field **2508** includes at least three sub-fields, including a hash identifier (HID) sub-field **2510**, a block length (BL) sub-field **2512**, and an extension (E) sub-field **2514**. HID sub-field **2510** indicates the type of hash function being applied to the message block. For example, in one implementation, HID sub-field **2510** may include two bits specifying one of the following hash types: SHA3, SHAKE, SHA2 (64-bit words), and SHA2 (32-bit words). BL sub-field **2512** indicates (possibly when interpreted together with HID sub-field **2510**) the length of the message block in bytes. E sub-field **2514** indicates whether or not the wide vector register **317** specified by register1 field **2504** holds the leading segment S0 or trailing segment S1 of the message block. For example, in an embodiment in which wide vector registers **317** are 1024 bits wide, E sub-field **2514** may have a value of b0 if the wide vector register **317** specified by register1 field **2504** does not hold the trailing segment of a message block and may have a value of b1 if the specified wide vector register **317** holds the trailing segment of the message block. Of course, in other embodiments in which wide vector registers **317** have different widths (e.g., 512 bits), E sub-field **2514** may include additional bits to specify additional register segments.

(114) Referring now to FIG. **26**, there is illustrated an exemplary padding circuit **2600** in accordance with one embodiment. Padding circuit **2600**, which can be implemented, for example, as a portion of data transfer circuit **406** of accelerator unit **314**, pads a message segment S1 held in a target wide vector register in response to execution of a padding instruction **2500** as shown in FIG. **25**. The illustrated example presumes a wide vector register file **316** having 1024-bit wide

vector registers **317**.

(115) In this exemplary embodiment, padding circuit **2600** includes select EOM circuit **2602**, which selects the value of the EOM padding byte **2302** or **2402** (i.e., eom_byte) based on the hash function specified by HID sub-field **2510** of padding instruction **2500**. Padding circuit **2600** also includes a select EOB circuit **2604** that, based on HID sub-field **2510**, similarly selects the value of the EOB padding byte (i.e., eob_byte) to be inserted by the padding instruction **2500**. In the described embodiment, for SHA3/SHAKE hash functions, select EOB circuit **2604** selects the fixed eob_byte value specified by the SHA-3 standard, which is contained in the register indicated by register2 field **2506**. For SHA2 hash functions, select EOB circuit **2604** selects a zero eob_byte because EOB padding words **2404** are inserted by a separate instruction in this embodiment.

Padding circuit **2600** further includes a select BL-size circuit **2606**, which selects and outputs an 8-bit block length value based on the HID field **2510** and BL field **2512** of padding instruction **2500**. (116) The 8-bit block length value output by select BL-size circuit **2606** is received by an EOB enable circuit **2608**, which includes a comparator **2610**, decoder **2612**, and bitwise AND circuit **2614**. The high order bit of the 8-bit block length value indicates whether or not the length of the message block exceeds the width of the 1024-bit wide vector register **317** (as would be the case, for example, for SHA3-224, SHAKE-128, and SHAKE 256). The lower order 7-bits of the block length form a block-length-size (bl_size) indicating the number of bytes comprising the segment of the message block buffered in the target wide vector register **317** identified by register1 field **2504**. Decoder **2612** decodes the 7-bit bl_size value to obtain a 128-bit representation of the location of the end of the message block within the target wide vector register **317**. Comparator **2610** compares the high order bit of the 8-bit block length to E sub-field **2514** of padding instruction **2500** to form a 1-bit indication of whether or not the EOB padding is to be added to the segment of the message block buffered in the target wide vector register (i.e., whether target wide vector register **317** buffers the trailing segment S1 of the message block). This 1-bit indication is then logically combined by bitwise AND circuit **2614** to produce a 128-bit EOB enable signal (eob_en(0:127)) identifying the byte (if any) of the message segment buffered in the target wide vector register **317** at which the EOB padding is to be inserted.

(117) Still referring to FIG. **26**, padding circuit **2600** further includes an EOM enable circuit **2620**, which includes a selection circuit **2620**, a comparator **2610**, decoder **2612**, and bitwise AND circuit **2614**. In the depicted example, selection circuit **2620** is implemented with a two-input multiplexer having a first input coupled to receive an 8-bit indication of message length and a second input coupled to receive an expanded message length applicable to SHA-2 hash functions employing 32-bit words. The expanded message length value at the second input doubles the original length of the message according to the equation $EX_LEN=4*(LEN/4)+LEN$ by the insertion of b0 between bits 5 and 6 of the original length. This technique preserves the bit locations of original bits 6:7, which indicate the byte location, if any, within the 32 high-order bits of the expanded message block of the final message byte. Selection circuit **2620** selects the first of its two 8-bit inputs if HID sub-field **2510** indicates the hash function is a SHA3/SHAKE hash function or a SHA2 hash function employing 64-bit words and alternatively selects the second of its two inputs if HID sub-field **2510** indicate the hash function is a SHA2 hash function employing 32-bit words.

(118) The 8-bit length value output by selection circuit **2620** includes a high order bit that indicates whether or not the block length exceeds the width of the 1024-bit wide vector register file **316** and seven lower order bits indicating the number of bytes comprising the segment of the message block buffered in the target wide vector register **317** identified by register1 field **2504**. Decoder **2612** decodes the seven low order bits to obtain a 128-bit representation of the byte location, if any, at which the end of the message byte is to be inserted within the target wide vector register **317**. Comparator **2622** compares the high order bit of the length value output by selection circuit **2620** to E sub-field **2514** of padding instruction **2500** to form a 1-bit indication of whether or not the EOM padding is to be added to the segment of the message block buffered in the target wide vector

register **317**. This 1-bit indication is then logically combined by bitwise AND circuit **2626** to produce a 128-bit EOM enable signal (eom_en(0:127)) identifying the byte (if any) of the message segment buffered in the target wide vector register **317** at which the EOM padding is to be inserted. (119) The EOB enable signal eob_en(0:127), the EOM enable signal eom_en(0:127), eom_byte, eob_byte, and the message segment from the target wide vector register **317** are all passed to a conditional OR circuit **2630**, which conditionally inserts the EOM and/or EOB padding into the message segment to obtain a padded message segment Sp. Padded message segment Sp is then stored back to the target wide vector register **317** specified in register1 field **2504**.

(120) With reference now to FIG. **27**, there is illustrated an exemplary embodiment of conditional OR circuit **2630** of FIG. **26**. In this example, each of the 128 bytes of message segment has a respective associated OR gate **2700** having three 8-bit inputs. The first input of OR gate **2700** is coupled to receive a respective byte of message segment S. The second input of OR gate **2700** is coupled to the output of a two-input AND gate **2702** that qualifies the eom_byte with the eom_en() for the given byte of message segment S. The third input of OR gate **2700** is coupled to the output of a two-input AND gate **2704** that qualifies the eob_byte with the eob_en() for the given byte of message segment S. OR gate **2700** performs a logical OR operation on these three inputs and writes the resulting byte of the padded message segment Sp to the target wide vector register **317** in wide vector register file **316**. Thus, if neither eom_en() nor eob_en() is asserted for a given byte of message segment S, the relevant OR gate **2700** simply writes the byte of input message segment S to the corresponding byte of padded message segment Sp. If, however, one or both of eom_en() or eob_en() is asserted for a given byte of message segment S, the relevant OR gate **2700** writes the eom_byte, the eob_byte, or the logical combination thereof into the corresponding byte of padded message segment Sp, as indicated by enable signals eom_en() and eob_en().

(121) Referring now to FIG. **28**, there is depicted a high-level logical flowchart of an exemplary process for padding a message block in accordance with one embodiment. The illustrated process can be performed by accelerator unit **314** in response to receipt of a padding instruction **2500**. For ease of understanding, the process is described below with reference to the exemplary padding circuit depicted in FIGS. **26-27**.

(122) The process of FIG. **28** begins at block **2800** and then proceeds to block **2802**, which illustrates accelerator unit **314** receiving a padding instruction **2500** for execution. In response to receipt of padding instruction **2500**, accelerator unit **314** first accesses the source operands specified by register fields **2504**, **2506** of padding instruction **2500** (block **2804**). In particular, accelerator unit **314** reads a message segment S from the target wide vector register **317** in wide vector register file **316** specified by register1 field **2504**, reads the unpadded message length from the register **301** in architected register file **300** specified by register2 field **2506**, and transfers these operands to padding circuit **2600** of FIG. **26**, which as noted above may be implemented within data transfer circuit **406**. At block **2806**, padding circuit **2600** utilizes the mode field **2508** of padding instruction **2500** to select parameters of the padding operation. In particular, select EOM circuit **2602** selects the value of the EOM padding byte (eom_byte) **2302** or **2402** based on the hash function specified by mode field **2508**, select EOB circuit **2604** selects the value of the EOB padding byte (eob_byte) to be inserted by the padding instruction **2500** (i.e., a fixed value for SHA3/SHAKE and a zero byte for SHA2 since the EOB padding words **2404** are applied by a separate instruction for SHA2), and select BL-size circuit **2606** selects the block length based on the HID sub-field **2510** and BL sub-field **2512**. The eom_byte selected by select EOM circuit **2602** and the eob_byte selected by selected EOB circuit **2604** form inputs to conditional OR circuit **2630**.

(123) At block **2808**, selection circuit **2620** determines based on the HID sub-field of mode field **2508** whether or not the hash function applied to the message is one of the SHA2-224 or SHA2-256 hash functions that employs 32-bit words. If not, selection circuit **2620** selects and outputs the message length read from the register **301** identified by register2 field **2506** as the length of the

message, and the process of FIG. 28 proceeds to block 2812, which is described below. If, however, selection circuit 2620 determines at block 2808 that HID sub-field 2510 of padding instruction 2500 indicates a SHA2 hash function employing 32-bit words, selection circuit 2620 selects and outputs a doubled length for the SHA2 message to account for the message expansion described above with reference to FIG. 15. In one implementation, the expanded SHA2 message length can conveniently be computed as: $4*(LEN/4)+LEN$. The process then proceeds from block 2810 to block 2812.

(124) Block 2812 illustrates a determination by EOM enable circuit 2620 whether or not EOM padding is to be placed in the current message segment. If not, the EOM enable vector `eom_en(0:127)` generated by EOM enable circuit 2620 is all zeros, and no EOM padding is inserted into message segment S. Accordingly, the process passes to block 2816, which is described below. If, however, EOM enable circuit 2620 determines at block 2812 that the EOM padding is to be inserted into message segment S, EOM enable circuit 2620 generates an EOM enable vector `eom_en(0:127)` identifying the byte of message segment S at which the EOM padding byte is to be inserted, and the EOM padding byte is inserted into the specified byte of padded message segment Sp by conditional OR circuit 2630 (block 2814). The process proceeds from block 2814 to block 2816.

(125) At block 2816, select BL-size circuit 2606 and EOB enable circuit 2608 determine whether or not the hash function specified by mode field 2508 of hash instruction 2500 is either a SHA3 or SHAKE hash function and an EOB padding byte is to be inserted in the message segment S. If not, the EOB enable vector `eob_en(0:127)` generated by EOB enable circuit 2608 is all zeros, and no EOB padding is inserted into message segment S. Accordingly, the process from block 2816 to passes to block 2820, which is described below. If, however, BL-size circuit 2606 and EOB enable circuit 2608 determine at block 2816 that the hash function specified by mode field 2508 is either a SHA3 or SHAKE hash function and EOB padding is to be inserted into message segment S, EOB enable circuit 2608 generates an EOB enable vector `eob_en(0:127)` identifying the byte of message segment S at which the EOB padding byte is to be inserted, and the EOB padding byte is inserted into the specified byte of padded message segment Sp by conditional OR circuit 2630 (block 2818). The process then passes to block 2820.

(126) Block 2820 illustrates data transfer circuit 406 writing the resultant padded message segment Sp into the target wide vector register 317 specified by register1 field 2504. Thereafter, the process of FIG. 28 ends at block 2822.

(127) With reference now to FIG. 29, there is illustrated a block diagram of an exemplary design flow 2900 used for example, in semiconductor IC logic design, simulation, test, layout, and manufacture. Design flow 2900 includes processes, machines and/or mechanisms for processing design structures or devices to generate logically or otherwise functionally equivalent representations of the design structures and/or devices described above and shown herein. The design structures processed and/or generated by design flow 2900 may be encoded on machine-readable transmission or storage media to include data and/or instructions that when executed or otherwise processed on a data processing system generate a logically, structurally, mechanically, or otherwise functionally equivalent representation of hardware components, circuits, devices, or systems. Machines include, but are not limited to, any machine used in an IC design process, such as designing, manufacturing, or simulating a circuit, component, device, or system. For example, machines may include: lithography machines, machines and/or equipment for generating masks (e.g. e-beam writers), computers or equipment for simulating design structures, any apparatus used in the manufacturing or test process, or any machines for programming functionally equivalent representations of the design structures into any medium (e.g. a machine for programming a programmable gate array).

(128) Design flow 2900 may vary depending on the type of representation being designed. For example, a design flow 2900 for building an application specific IC (ASIC) may differ from a

design flow **2900** for designing a standard component or from a design flow **2900** for instantiating the design into a programmable array, for example a programmable gate array (PGA) or a field programmable gate array (FPGA) offered by Altera® Inc. or Xilinx® Inc.

(129) FIG. **29** illustrates multiple such design structures including an input design structure **1020** that is preferably processed by a design process **2910**. Design structure **2920** may be a logical simulation design structure generated and processed by design process **2910** to produce a logically equivalent functional representation of a hardware device. Design structure **2920** may also or alternatively comprise data and/or program instructions that when processed by design process **2910**, generate a functional representation of the physical structure of a hardware device. Whether representing functional and/or structural design features, design structure **2920** may be generated using electronic computer-aided design (ECAD) such as implemented by a core developer/designer. When encoded on a machine-readable data transmission, gate array, or storage medium, design structure **2920** may be accessed and processed by one or more hardware and/or software modules within design process **2910** to simulate or otherwise functionally represent an electronic component, circuit, electronic or logic module, apparatus, device, or system such as those shown herein. As such, design structure **2920** may comprise files or other data structures including human and/or machine-readable source code, compiled structures, and computer-executable code structures that when processed by a design or simulation data processing system, functionally simulate or otherwise represent circuits or other levels of hardware logic design. Such data structures may include hardware-description language (HDL) design entities or other data structures conforming to and/or compatible with lower-level HDL design languages such as Verilog and VHDL, and/or higher-level design languages such as C or C++.

(130) Design process **2910** preferably employs and incorporates hardware and/or software modules for synthesizing, translating, or otherwise processing a design/simulation functional equivalent of the components, circuits, devices, or logic structures shown herein to generate a netlist **2980** which may contain design structures such as design structure **2920**. Netlist **2980** may comprise, for example, compiled or otherwise processed data structures representing a list of wires, discrete components, logic gates, control circuits, I/O devices, models, etc. that describes the connections to other elements and circuits in an integrated circuit design. Netlist **2980** may be synthesized using an iterative process in which netlist **2980** is resynthesized one or more times depending on design specifications and parameters for the device. As with other design structure types described herein, netlist **2980** may be recorded on a machine-readable storage medium or programmed into a programmable gate array. The medium may be a non-volatile storage medium such as a magnetic or optical disk drive, a programmable gate array, a compact flash, or other flash memory.

Additionally, or in the alternative, the medium may be a system or cache memory, or buffer space.

(131) Design process **2910** may include hardware and software modules for processing a variety of input data structure types including netlist **2980**. Such data structure types may reside, for example, within library elements **2930** and include a set of commonly used elements, circuits, and devices, including models, layouts, and symbolic representations, for a given manufacturing technology (e.g., different technology nodes, 32 nm, 45 nm, 290 nm, etc.). The data structure types may further include design specifications **2940**, characterization data **2950**, verification data **2960**, design rules **2990**, and test data files **2985** which may include input test patterns, output test results, and other testing information. Design process **2910** may further include, for example, standard mechanical design processes such as stress analysis, thermal analysis, mechanical event simulation, process simulation for operations such as casting, molding, and die press forming, etc. One of ordinary skill in the art of mechanical design can appreciate the extent of possible mechanical design tools and applications used in design process **2910** without deviating from the scope and spirit of the invention. Design process **2910** may also include modules for performing standard circuit design processes such as timing analysis, verification, design rule checking, place and route operations, etc.

(132) Design process **2910** employs and incorporates logic and physical design tools such as HDL compilers and simulation model build tools to process design structure **2920** together with some or all of the depicted supporting data structures along with any additional mechanical design or data (if applicable), to generate a second design structure **2990**. Design structure **2990** resides on a storage medium or programmable gate array in a data format used for the exchange of data of mechanical devices and structures (e.g., information stored in a IGES, DXF, Parasolid XT, JT, DRG, or any other suitable format for storing or rendering such mechanical design structures). Similar to design structure **2920**, design structure **2990** preferably comprises one or more files, data structures, or other computer-encoded data or instructions that reside on transmission or data storage media and that when processed by an ECAD system generate a logically or otherwise functionally equivalent form of one or more of the embodiments of the invention shown herein. In one embodiment, design structure **2990** may comprise a compiled, executable HDL simulation model that functionally simulates the devices shown herein.

(133) Design structure **2990** may also employ a data format used for the exchange of layout data of integrated circuits and/or symbolic data format (e.g., information stored in a GDSII (GDS2), GL1, OASIS, map files, or any other suitable format for storing such design data structures). Design structure **2990** may comprise information such as, for example, symbolic data, map files, test data files, design content files, manufacturing data, layout parameters, wires, levels of metal, vias, shapes, data for routing through the manufacturing line, and any other data required by a manufacturer or other designer/developer to produce a device or structure as described above and shown herein. Design structure **2990** may then proceed to a stage **2995** where, for example, design structure **2990**: proceeds to tape-out, is released to manufacturing, is released to a mask house, is sent to another design house, is sent back to the customer, etc.

(134) As has been described, in at least one embodiment, a processor includes a register file and an execution unit. The execution unit includes a hash circuit including at least a state register, a state update circuit coupled to the state register, and a control circuit. Based on a hash instruction, the hash circuit receives from the register file and buffers within the state register a current state of a message being hashed. The state update circuit performs state update function on contents of the state register, where performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register. Following completion of all of the plurality of iterative rounds of processing, the execution unit stores contents of the state register to the register file as an updated state of the message.

(135) In at least some embodiments, the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function and the state update circuit performs twenty-four rounds of processing each utilizes as an input a respective one of twenty-four round indices.

(136) In one embodiment, the execution unit executes the hash instruction in a squeezing phase of a Secure Hash Algorithm and Keccak (SHAKE) hash algorithm.

(137) In at least some embodiments, the state update function comprises a Secure Hash Algorithm 2 (SHA2) block hash function. In at least some embodiments, the hash circuit further includes an adder configured to add contents of the state register to the current state and return a resulting sum to the register file. In some embodiments, the execution unit further includes a message block register for buffering a message block of the message and a message schedule round circuit coupled to the message block register. The message schedule round circuit performs a plurality of iterative rounds of processing on contents of the message block register and returns a result of each of the plurality of iterative rounds of processing to the message block register. In some embodiments, the state update circuit includes a data path for data words having a first data width and the execution unit is configured, based on the hash instruction indicating a second data width that is narrower than the first data width, to expand data words of a message block of the message to the first data width prior to processing the data words of the message in the state update circuit.

(138) While various embodiments have been particularly shown and described, it will be understood by those skilled in the art that various changes in form and detail may be made therein without departing from the spirit and scope of the appended claims and these alternate implementations all fall within the scope of the appended claims. For example, while the invention has been described with specific reference to the SHA family of standards, those skilled in the art will appreciate that the disclosed inventions are also applicable to other hash algorithms (e.g., general Keccak functions, among others). Further, although illustrative numbers of bits and bytes have been discussed herein for ease of understanding, it should be appreciated that the specific numbers of bits and bytes employed in hashing algorithms can and do change over time and that the principles of the disclosed inventions are applicable to cryptographic algorithms regardless of the specific numbers of bits and bytes in a given implementation.

(139) The flowchart and block diagrams in the Figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods, and computer program products according to various embodiments of the present invention. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of instructions, which comprises one or more executable instructions for implementing the specified logical function(s). In some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts or carry out combinations of special purpose hardware and computer instructions.

(140) Further, although aspects have been described with respect to a computer system executing program code that directs the functions of the present invention, it should be understood that present invention may alternatively be implemented as a program product including a computer-readable storage device storing program code that can be processed by a data processing system. The computer-readable storage device can include volatile or non-volatile memory, an optical or magnetic disk, or the like. However, as employed herein, a “storage device” is specifically defined to include only statutory articles of manufacture and to exclude signal media per se, transitory propagating signals per se, and energy per se.

(141) The program product may include data and/or instructions that when executed or otherwise processed on a data processing system generate a logically, structurally, or otherwise functionally equivalent representation (including a simulation model) of hardware components, circuits, devices, or systems disclosed herein. Such data and/or instructions may include hardware-description language (HDL) design entities or other data structures conforming to and/or compatible with lower-level HDL design languages such as Verilog and VHDL, and/or higher-level design languages such as C or C++. Furthermore, the data and/or instructions may also employ a data format used for the exchange of layout data of integrated circuits and/or symbolic data format (e.g., information stored in a GDSII (GDS2), GL1, OASIS, map files, or any other suitable format for storing such design data structures).

Claims

1. A processor, comprising: an instruction fetch unit that fetches instructions to be processed in the processor; a register file including a plurality of registers for storing source operands and destination operands; and an execution unit for executing a hash instruction obtained from the instructions, wherein the execution unit includes a hash circuit including at least: a state register, a state update circuit coupled to the state register, a control circuit, and a single instruction multiple

data (SIMD) adder, wherein the SIMD adder is a dedicated component of the hash circuit, wherein the execution unit, based on the hash instruction, is configured to perform: receiving from the register file and buffering within the state register a current state of a message being hashed; performing, in the state update circuit, a state update function on contents of the state register, wherein: the state update function comprises a Secure Hash Algorithm 2 (SHA2) block hash function; performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register; and following completion of all of the plurality of iterative rounds of processing, adding, by the SIMD adder, contents of the state register to a current state of the message being hashed and storing a resulting sum to the register file as an updated state of the message.

2. The processor of claim 1, wherein: the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function; and performing a plurality of iterative rounds of processing includes performing twenty-four rounds of processing in which each round utilizes as an input a respective one of twenty-four round indices.

3. The processor of claim 1, wherein the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function; and performing a plurality of iterative rounds of processing includes performing a number of rounds of processing in indicated by the hash instruction.

4. The processor of claim 1, wherein the execution unit executes the hash instruction in a squeezing phase of a Secure Hash Algorithm and Keccak (SHAKE) hash algorithm.

5. The processor of claim 1, wherein: the execution unit further includes a message block register for buffering a message block of the message and a message schedule round circuit coupled to the message block register; and performing a state update function includes performing, by the message schedule round circuit, a plurality of iterative rounds of processing on contents of the message block register and returning a result of each of the plurality of iterative rounds of processing to the message block register.

6. The processor of claim 1, wherein: the state update circuit includes a data path for data words having a first data width; and the execution unit is configured, based on the hash instruction indicating a second data width that is narrower than the first data width, to expand data words of a message block of the message to the first data width prior to processing the data words of the message in the state update circuit.

7. A method of data processing in a processor, said method comprising: fetching, by an instruction fetch unit, instructions to be processed by the processor, wherein the instructions include a hash instruction; and based on receiving the hash instruction, an execution unit of the processor executing the hash instruction, wherein the execution unit includes a hash circuit including: at least a state register, a state update circuit coupled to the state register, a control circuit, and a single instruction multiple data (SIMD) adder, wherein the SIMD adder is a distinct pipeline shared by the hash circuit and at least one other hash circuit, wherein the executing includes: receiving from a register file and buffering within the state register a current state of a message being hashed; performing, in the state update circuit, a state update function on contents of the state register, wherein: the state update function comprises a Secure Hash Algorithm 2 (SHA2) block hash function; performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register; and following completion of all of the plurality of iterative rounds of processing, adding, by the SIMD adder, contents of the state register to a current state of a message being hashed and storing a resulting sum to the register file as an updated state of the message.

8. The method of claim 7, wherein: the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function; and performing a plurality of iterative rounds of processing includes performing twenty-four rounds of processing in which each round utilizes as an input a

respective one of twenty-four round indices.

9. The method of claim 7, wherein executing the hash instruction includes executing the hash instruction in a squeezing phase of a Secure Hash Algorithm and Keccak (SHAKE) hash algorithm.

10. The method of claim 7, wherein: the execution unit further includes a message block register for buffering a message block of the message and a message schedule round circuit coupled to the message block register; and performing a state update function includes performing a plurality of iterative rounds of processing on contents of the message block register in the message schedule round circuit and returning a result of each of the plurality of iterative rounds of processing to the message block register.

11. The method of claim 7, wherein: the state update circuit includes a data path for data words having a first data width; and the method further comprises: based on the hash instruction indicating a second data width that is narrower than the first data width, expanding data words of a message block of the message to the first data width prior to processing the data words of the message in the state update circuit.

12. A program product, comprising: a machine-readable storage device; and a design structure embodied in the machine-readable storage device, wherein the design structure is for designing, manufacturing, or testing an integrated circuit, the design structure including: a processor, including: an instruction fetch unit that fetches instructions to be executed; a register file including a plurality of registers for storing source and destination operands; and an execution unit for executing a hash instruction obtained from the instructions, wherein the execution unit includes a hash circuit including at least: a state register, a state update circuit coupled to the state register, a control circuit, and a single instruction multiple data (SIMD) an adder, wherein the SIMD adder is a distinct pipeline shared by the hash circuit and at least one other hash circuit, wherein the execution unit, based on the hash instruction, is configured to perform: receiving from the register file and buffering within the state register a current state of a message being hashed; performing, in the state update circuit, a state update function on contents of the state register, wherein: the state update function comprises a Secure Hash Algorithm 2 (SHA2) block hash function; performing the state update function includes performing a plurality of iterative rounds of processing on contents of the state register and returning a result of each of the plurality of iterative rounds of processing to the state register; and following completion of all of the plurality of iterative rounds of processing, adding by the SIMD adder, contents of the state register to a current state of a message being hashed and storing a resulting sum to the register file as an updated state of the message.

13. The program product of claim 12, wherein: the state update function comprises a Secure Hash Algorithm 3 (SHA3) state permute function; and performing a plurality of iterative rounds of processing includes performing twenty-four rounds of processing in which each round utilizes as an input a respective one of twenty-four round indices.

14. The program product of claim 12, wherein the execution unit executes the hash instruction in a squeezing phase of a Secure Hash Algorithm and Keccak (SHAKE) hash algorithm.

15. The program product of claim 12, wherein: the execution unit further includes a message block register for buffering a message block of the message and a message schedule round circuit coupled to the message block register; and performing a state update function includes performing, by the message schedule round circuit, a plurality of iterative rounds of processing on contents of the message block register and returning a result of each of the plurality of iterative rounds of processing to the message block register.

16. The program product of claim 12, wherein: the state update circuit includes a data path for data words having a first data width; and the execution unit is configured, based on the hash instruction indicating a second data width that is narrower than the first data width, to expand data words of a message block of the message to the first data width prior to processing the data words of the message in the state update circuit.
