

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12412072
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Markram; Henry et al.

Characterizing activity in a recurrent artificial neural network

Abstract

Methods, systems, and apparatus, including computer programs encoded on a computer storage medium, for characterizing activity in a recurrent artificial neural network. In one aspect, a method includes outputting digits from a recurrent artificial neural network, wherein each digit represents whether or not activity within a particular group of nodes in the recurrent artificial neural network comports with a respective pattern of activity.

Inventors: Markram; Henry (Lausanne, CH), Levi; Ran (Aberdeen, GB), Bellwald; Kathryn Pamela Hess (Aigle, CH)

Applicant: INAIT SA (Lausanne, CH)

Family ID: 68765152

Assignee: INAIT SA (Lausanne, CH)

Appl. No.: 16/004796

Filed: June 11, 2018

Prior Publication Data

Document Identifier	Publication Date
US 20190378000 A1	Dec. 12, 2019

Publication Classification

Int. Cl.: G06N3/044 (20230101); G06N3/08 (20230101)

U.S. Cl.:

CPC G06N3/044 (20230101); G06N3/08 (20130101);

Field of Classification Search

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5822742	12/1997	Alkon et al.	N/A	N/A
7321882	12/2007	Herbert	N/A	N/A
7412426	12/2007	Hercus	N/A	N/A
8818923	12/2013	Hoffmann	N/A	N/A
9336239	12/2015	Hoffmann et al.	N/A	N/A
9558442	12/2016	Canoy et al.	N/A	N/A
9785886	12/2016	Andoni et al.	N/A	N/A
9875440	12/2017	Lamport	N/A	N/A
10019506	12/2017	Li et al.	N/A	N/A
10153806	12/2017	Petre et al.	N/A	N/A
10417558	12/2018	Bauer et al.	N/A	N/A
10510000	12/2018	Lamport	N/A	N/A
10621364	12/2019	El Defrawy	N/A	N/A
10628486	12/2019	Chu et al.	N/A	N/A
10650047	12/2019	Yanagisawa	N/A	N/A
10885020	12/2020	Ablitt	N/A	N/A
10922510	12/2020	Tscherepanow et al.	N/A	N/A
11195038	12/2020	Nunn et al.	N/A	N/A
11218498	12/2021	Hajimirsadeghi et al.	N/A	N/A
11250326	12/2021	Ko et al.	N/A	N/A
11278413	12/2021	Lang	N/A	N/A
11569978	12/2022	Hess et al.	N/A	N/A
11580401	12/2022	Markram et al.	N/A	N/A
11615285	12/2022	Reimann et al.	N/A	N/A
11651210	12/2022	Henry et al.	N/A	N/A
11652603	12/2022	Markram et al.	N/A	N/A
11663478	12/2022	Markram et al.	N/A	N/A
11797827	12/2022	Markram et al.	N/A	N/A
11816553	12/2022	Markram et al.	N/A	N/A
11893471	12/2023	Markram et al.	N/A	N/A
11972343	12/2023	Markram et al.	N/A	N/A
12020157	12/2023	Markram et al.	N/A	N/A
12113891	12/2023	Hess Bellwald et al.	N/A	N/A
12147904	12/2023	Markram et al.	N/A	N/A
12154023	12/2023	Markram et al.	N/A	N/A
2003/0140020	12/2002	Chen et al.	N/A	N/A
2004/0015459	12/2003	Herbert	N/A	N/A
2006/0112028	12/2005	Xiao et al.	N/A	N/A
2006/0184471	12/2005	Minamino et al.	N/A	N/A
2008/0152217	12/2007	Greer	N/A	N/A
2008/0267188	12/2007	Furber et al.	N/A	N/A
2009/0006060	12/2008	Rhodes	703/11	G06N 3/047
2009/0012581	12/2008	Rhodes	N/A	N/A

2009/0187736	12/2008	Raichelgauz et al.	N/A	N/A
2014/0156901	12/2013	Raichelgauz et al.	N/A	N/A
2015/0058352	12/2014	Brand	N/A	N/A
2015/0206049	12/2014	Canoy et al.	N/A	N/A
2015/0280906	12/2014	Shany et al.	N/A	N/A
2015/0310303	12/2014	Andreopoulos et al.	N/A	N/A
2015/0347870	12/2014	Andreopoulos et al.	N/A	N/A
2015/0363689	12/2014	Henry et al.	N/A	N/A
2015/0371132	12/2014	Gemello et al.	N/A	N/A
2016/0048756	12/2015	Hall et al.	N/A	N/A
2016/0350648	12/2015	Gilad-Bachrach et al.	N/A	N/A
2017/0024661	12/2016	Hunsberger et al.	N/A	N/A
2017/0124452	12/2016	Tucker et al.	N/A	N/A
2017/0139759	12/2016	Bandara	N/A	G06F 11/0706
2017/0213131	12/2016	Hammond et al.	N/A	N/A
2017/0229117	12/2016	van der Made et al.	N/A	N/A
2017/0249536	12/2016	Hillar et al.	N/A	N/A
2018/0005111	12/2017	Chaudhari et al.	N/A	N/A
2018/0018553	12/2017	Bach et al.	N/A	N/A
2018/0053108	12/2017	Olabiya et al.	N/A	N/A
2018/0053114	12/2017	Adjaoute	N/A	N/A
2018/0082171	12/2017	Merity et al.	N/A	N/A
2018/0107922	12/2017	Paul et al.	N/A	N/A
2018/0136912	12/2017	Venkataramani et al.	N/A	N/A
2018/0157973	12/2017	El-Yaniv et al.	N/A	N/A
2018/0165547	12/2017	Huang et al.	N/A	N/A
2018/0197069	12/2017	Reimann et al.	N/A	N/A
2018/0197076	12/2017	Paik et al.	N/A	N/A
2018/0247198	12/2017	Vasudevan et al.	N/A	N/A
2018/0336453	12/2017	Merity et al.	N/A	N/A
2018/0373985	12/2017	Yang et al.	N/A	N/A
2019/0012574	12/2018	Anthony et al.	N/A	N/A
2019/0036678	12/2018	Ahmed	N/A	N/A
2019/0042988	12/2018	Brown et al.	N/A	N/A
2019/0044916	12/2018	Jones	N/A	N/A
2019/0050726	12/2018	Azaria et al.	N/A	N/A
2019/0122096	12/2018	Husain	N/A	N/A
2019/0122140	12/2018	Sen	N/A	N/A
2019/0171187	12/2018	Cella et al.	N/A	N/A
2019/0171929	12/2018	Abadi et al.	N/A	N/A
2019/0180187	12/2018	Rawal et al.	N/A	N/A
2019/0197410	12/2018	Berry, II	N/A	N/A
2019/0228300	12/2018	Cao et al.	N/A	N/A
2019/0244348	12/2018	Buckler et al.	N/A	N/A
2019/0265950	12/2018	Heeger et al.	N/A	N/A
2019/0266246	12/2018	Wang et al.	N/A	N/A
2019/0286074	12/2018	Hoffmann	N/A	N/A
2019/0304568	12/2018	Wei et al.	N/A	N/A
2019/0312898	12/2018	Verma et al.	N/A	N/A
2019/0335192	12/2018	Otto	N/A	H04N 19/137

2019/0354846	12/2018	Mellempudi et al.	N/A	N/A
2019/0370647	12/2018	Doshi et al.	N/A	N/A
2019/0377976	12/2018	Markram et al.	N/A	N/A
2019/0377999	12/2018	Markram et al.	N/A	N/A
2019/0378007	12/2018	Markram et al.	N/A	N/A
2019/0378008	12/2018	Markram et al.	N/A	N/A
2019/0392303	12/2018	Cherubini et al.	N/A	N/A
2020/0011668	12/2019	Derhy et al.	N/A	N/A
2020/0012927	12/2019	Raichelgauz et al.	N/A	N/A
2020/0098135	12/2019	Ganjineh et al.	N/A	N/A
2020/0117690	12/2019	Tran	N/A	N/A
2020/0126126	12/2019	Briancon et al.	N/A	N/A
2020/0134016	12/2019	Cao et al.	N/A	N/A
2020/0184055	12/2019	Storm et al.	N/A	N/A
2020/0210764	12/2019	Hamedi et al.	N/A	N/A
2020/0228336	12/2019	Streit	N/A	N/A
2020/0242444	12/2019	Zhang et al.	N/A	N/A
2020/0285944	12/2019	Lee et al.	N/A	N/A
2020/0302297	12/2019	Jaganathan et al.	N/A	N/A
2020/0304284	12/2019	Markram et al.	N/A	N/A
2020/0304285	12/2019	Hess et al.	N/A	N/A
2020/0310400	12/2019	Jha et al.	N/A	N/A
2020/0367810	12/2019	Shouldice et al.	N/A	N/A
2020/0380335	12/2019	Neznal	N/A	N/A
2020/0402497	12/2019	Semonov et al.	N/A	N/A
2021/0027163	12/2020	Baker	N/A	N/A
2021/0049441	12/2020	Bronstein	N/A	N/A
2021/0049446	12/2020	Gurumurthi et al.	N/A	N/A
2021/0058547	12/2020	Puttamalla et al.	N/A	N/A
2021/0097578	12/2020	Holmes et al.	N/A	N/A
2021/0110115	12/2020	Hermann et al.	N/A	N/A
2021/0182604	12/2020	Anthony et al.	N/A	N/A
2021/0182653	12/2020	Markram et al.	N/A	N/A
2021/0182654	12/2020	Markram et al.	N/A	N/A
2021/0182655	12/2020	Markram et al.	N/A	N/A
2021/0182657	12/2020	Markram et al.	N/A	N/A
2021/0182681	12/2020	Markram et al.	N/A	N/A
2021/0271319	12/2020	Lussier et al.	N/A	N/A
2021/0319098	12/2020	Pogorelik et al.	N/A	N/A
2021/0338007	12/2020	Choi et al.	N/A	N/A
2021/0398621	12/2020	Stojevic et al.	N/A	N/A
2022/0005332	12/2021	Metzler et al.	N/A	N/A
2022/0012877	12/2021	Buckler et al.	N/A	N/A
2022/0147760	12/2021	Dutta et al.	N/A	N/A
2022/0148454	12/2021	Jaramaz et al.	N/A	N/A
2022/0157436	12/2021	Harley et al.	N/A	N/A
2022/0187847	12/2021	Cella et al.	N/A	N/A
2022/0197306	12/2021	Cella et al.	N/A	N/A
2022/0261593	12/2021	Yu et al.	N/A	N/A
2022/0269346	12/2021	Hussami et al.	N/A	N/A

2023/0019839	12/2022	Markram	N/A	N/A
2023/0024925	12/2022	Markram et al.	N/A	N/A
2023/0028511	12/2022	Markram et al.	N/A	N/A
2023/0085384	12/2022	Lütgehetmann et al.	N/A	N/A
2023/0171086	12/2022	Hess et al.	N/A	N/A
2023/0297808	12/2022	Reimann et al.	N/A	N/A
2023/0316077	12/2022	Henry et al.	N/A	N/A
2023/0351196	12/2022	Markram et al.	N/A	N/A
2023/0370244	12/2022	Markram et al.	N/A	N/A
2024/0046077	12/2023	Markram et al.	N/A	N/A
2024/0111994	12/2023	Markram et al.	N/A	N/A
2024/0176985	12/2023	Markram et al.	N/A	N/A
2024/0362480	12/2023	Markram et al.	N/A	N/A
2024/0386265	12/2023	Markram et al.	N/A	N/A
2025/0070956	12/2024	Hess Bellwald et al.	N/A	N/A
2025/0117628	12/2024	Markram et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
101669130	12/2009	CN	N/A
101719908	12/2009	CN	N/A
102859538	12/2012	CN	N/A
103988203	12/2013	CN	N/A
104318304	12/2014	CN	N/A
104335219	12/2014	CN	N/A
105095966	12/2014	CN	N/A
105787561	12/2015	CN	N/A
106156845	12/2015	CN	N/A
107247989	12/2016	CN	N/A
107423814	12/2016	CN	N/A
107743630	12/2017	CN	N/A
107844830	12/2017	CN	N/A
108348750	12/2017	CN	N/A
112567388	12/2020	CN	N/A
112567389	12/2020	CN	N/A
1283496	12/2002	EP	N/A
3340121	12/2017	EP	N/A
3803707	12/2020	EP	N/A
3803708	12/2020	EP	N/A
20210008418	12/2020	KR	N/A
20210008419	12/2020	KR	N/A
201437945	12/2013	TW	N/A
201535277	12/2014	TW	N/A
201725519	12/2016	TW	N/A
I608429	12/2016	TW	N/A
201928789	12/2018	TW	N/A
201935326	12/2018	TW	N/A
201937392	12/2018	TW	N/A
WO 2007/137047	12/2006	WO	N/A
WO 2016/206765	12/2015	WO	N/A
WO 2017/083399	12/2016	WO	N/A

WO 2017/197375	12/2016	WO	N/A
WO 2017/200971	12/2016	WO	N/A
WO 2018/058426	12/2017	WO	N/A
WO 2018/175400	12/2017	WO	N/A
WO 2019/238522	12/2018	WO	N/A
WO 2019/238523	12/2018	WO	N/A
WO 2020/187676	12/2019	WO	N/A
WO 2020/187694	12/2019	WO	N/A
WO 2021/116071	12/2020	WO	N/A
WO 2021/116075	12/2020	WO	N/A
WO 2021/116140	12/2020	WO	N/A
WO 2021/116147	12/2020	WO	N/A
WO 2021/116250	12/2020	WO	N/A
WO 2021/116379	12/2020	WO	N/A
WO 2021/116402	12/2020	WO	N/A
WO 2021/116404	12/2020	WO	N/A
WO 2021/116407	12/2020	WO	N/A

OTHER PUBLICATIONS

Blanco et al. "A genetic algorithm to obtain the optimal recurrent neural network." International Journal of Approximate Reasoning 23 (2000) 67-83 (Year: 2000). cited by examiner

Lee et al. "Detecting highly overlapping community structure by greedy clique expansion," arXiv preprint (arXiv:1002.1827v2), 2010, 10 pages (Year: 2010). cited by examiner

Dlotko et al. "Topological analysis of the connectome of digital reconstructions of neural microcircuits," arXiv preprint, 2016, 28 pages (IDS Oct. 4, 2019, doc. 2) (Year: 2016). cited by examiner

Zhang et al., "Ternary weight networks," arXiv preprint (arXiv:1605.04711v2), 2016, 5 pages (Year: 2016). cited by examiner

Masulli et al. The topology of the directed clique complex as a network invariant. Springerplus. 2016;5:388. Published Mar. 31, 2016. 12 pages. (Year: 2016). cited by examiner

Paweł Dłotko, "Directed clique topology" (slides are titled "Directed Complexes"), May 11, 2017, BIRS, Topological Methods in Brain Network Analysis (17w5043), Accessed from <http://www.birs.ca/events/2017/5-day-workshops/17w5043/videos/watch/201705111127-Dlotko.html> (Year: 2017). cited by examiner

TW Office Action in Taiwanese Appln. No. 108119813, dated Jun. 29, 2020, 17 pages (with machine translation). cited by applicant

Dlotko et al., "Topological analysis of the connectome of digital reconstructions of neural microcircuits," arXiv preprint, 2016, 28 pages. cited by applicant

Paugam-Moisy et al, "Delay learning and polychronization for reservoir computing," Neurocomputing, Mar. 2008, 71(7-9):1143-1158. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2019/064593, dated Sep. 6, 2019, 18 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2019/064740, dated Sep. 6, 2019, 18 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2019/064741, dated Sep. 6, 2019, 16 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2019/064773, dated Sep. 6, 2019, 17 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2019/064776, dated Sep. 6, 2019, 17 pages. cited by applicant

Reimann et al, "Cliques of Neurons Bound into Cavities Provide a Missing Link between Structure and

Function,” *Frontiers in Computational Neuroscience*, Jun. 2017, 11:48:16 pages. cited by applicant

Sun et al, “Runtime detection of activated polychronous neuronal group towards its spatiotemporal analysis,” 2015 International Joint Conference on Neural Networks, Jul. 2015, 1-8. cited by applicant

Woodward & Ikegami, “A Reservoir Computing approach to Image Classification using Coupled Echo State and Back-Propagation Neural Networks,” International conference image and vision computing, Auckland, New Zealand, Nov. 29, 2011, 6 pages. cited by applicant

PCT International Preliminary Report on Patentability in International Appln. No. PCT/EP2019/064741, dated Dec. 24, 2020, 9 pages. cited by applicant

PCT International Preliminary Report on Patentability in International Appln. No. PCT/EP2019/064773, dated Dec. 24, 2020, 11 pages. cited by applicant

PCT International Preliminary Report on Patentability in International Appln. No. PCT/EP2019/064776, dated Dec. 24, 2020, 11 pages. cited by applicant

PCT International Preliminary Report on Patentability in International Appln. No. PCT/EP2019/064740, dated Dec. 24, 2020, 13 pages. cited by applicant

PCT International Preliminary Report on Patentability in International Appln. No. PCT/EP2019/064593, dated Dec. 24, 2020, 11 pages. cited by applicant

Abbas et al., “Artificial Intelligence Conquering the next frontier of the digital world,” Research Gate, Dec. 17, 2017, 15 pages. cited by applicant

Aharoni et al. “Eigenvalues and homology of flag complexes and vector representations of graphs,” *Geom. Funct. Anal.*, Jul. 28, 2005, 15:555-566. cited by applicant

Angeli, “Symmetric functions for fast image retrieval with persistent homology,” *Math Meth Appl Sci.*, Apr. 24, 2018, 41:9567-9577. cited by applicant

Bale et al., “Efficient population coding of naturalistic whisker motion in the ventro-posterior medial thalamus based on precise spike timing,” *Front. Neural Circuits*, Sep. 25, 2015, 9: 1-14. cited by applicant

Bassett et al., “Network Neuroscience,” *Network Neuroscience*, *Nat. Neurosci.*, Mar. 20, 2017, 20:353-364. cited by applicant

Bauer et al., “PHAT—Persistent Homology Algorithms Toolbox,” *J. Symb. Comput.*, Jan. 1, 2017, 78: 76-90. cited by applicant

Bergomi et al., “Topological Graph persistence,” *Mathematics*, Jul. 20, 2017, pp. 1-15. cited by applicant

Bienenstock et al., “A model of neocortex,” *Network Comput. Neural Syst.*, Jul. 27, 1994, 6:179-224. cited by applicant

Brody et al., “Correlations without synchrony,” *Neural Comput.* 11, Oct. 1, 1999, 11:1537-1551. cited by applicant

Bullmore et al., “Complex brain networks: graph theoretical analysis of structural and functional systems,” *Nat. Rev. Neurosci.*, Mar. 10, 2009, 10:186-198. cited by applicant

Chella et al., “Representational Issues in the Debate on the Standard Model of the Mind,” *Association for the Advancement of Artificial Intelligence*, Nov. 12, 2017, pp. 1-6. cited by applicant

Chklovskii et al., “Cortical rewiring and information storage,” *Nature*, Oct. 14, 2004, 431:782-788. cited by applicant

Choi et al., “A Survey of Binary Similarity and distance measures,” *Systemic, Cybernetics, and Informatics*, Jan. 8, 2010, 8:43-48. cited by applicant

Choudhary, “Approximation algorithms for Vietoris-Rips and Čech filtrations,” *Doctoral Thesis of Computer Science, Faculty of Mathematics and Computer Science, Heidelberg University*, Dec. 14, 2017, 140 pages. cited by applicant

Courtney et al., “Dense Power-law Networks and Simplicial Complexes,” *ArXiv*, Apr. 23, 2018, pp. 1-16. cited by applicant

Crawford et al., “A theory on the role of π -electrons of docosahexaenoic acid in brain function: The six methylene-interrupted double bonds and the precision of neural signaling,” *Oilseeds & fats Crops and Lipids*, May 21, 2018, pp. 1-14. cited by applicant

Davydov et al., “Neural Network Structures: Current and Future States,” Research Gate, Feb. 2, 2018, pp. 1-6. cited by applicant

Engel et al., “Temporal binding and the neural correlates of sensory awareness,” Trends Cogn. Sci., Jan. 1, 2001, 5: 16-25. cited by applicant

Erdős et al., “On the evolution of random graphs,” Publ. Math. Inst. Hung. Acad. Sci., Jan. 1960, 17-60. cited by applicant

Github.com [online] “pytorch-tutorial/tutorials/03-advanced/image_captioning/,” Oct. 2007, retrieved on Jan. 21, 2020, retrieved from UR: <https://github.com/yunjey/pytorch-tutorial/tree/master/tutorials/03-advanced/image_captioning>, 3 pages. cited by applicant

Giusti et al., “Clique topology reveals intrinsic geometric structure in neural correlations,” Proc. Natl. Acad. Sci. U.S.A., Apr. 28, 2015, 112: 13455-13460. cited by applicant

Gong, “Evolution to a small-world network with chaotic units,” Europhysics Letters (EPL), Jul. 15, 2004, 67:328-333. cited by applicant

Harris et al. “The neocortical circuit: themes and variations,” Nat. Neurosci., Jan. 27, 2015, 18:170-181. cited by applicant

Hastings et al., “Challenges in the analysis of complex systems: introduction and overview,” The European Physical Journal Special Topics, Dec. 28, 2017, 226:3185-3197. cited by applicant

Haun et al., “Conscious Perception as Integrated Information Patterns in Human Electroencephalography,” eNeuro: Cognition and Behavior, Sep. 19, 2017, 4:2-18. cited by applicant

Hebb, “The Organization of Behavior: A Neuropsychological Theory,” New York, NY: Wiley & Sons, 1949, pp. 1-365. cited by applicant

Holtmaat et al., “Experience-dependent structural synaptic plasticity in the mammalian brain,” Nat. Rev. Neurosci., Sep. 10, 2009, 10: 647-658. cited by applicant

Ivancevic et al., “Tensor-Centric Warfare II: Entropic Uncertainty Modeling,” Intelligent Control and Automation, May 30, 2018, 9:30-51. cited by applicant

Judge, “Prefix “Re-Cognition” as prelude to fixing sustainability—“pro” vs. “con”? Speculative review of missing emphasis potentially vital for psychosocial balance,” Research Gate, Jul. 16, 2017, pp. 1-22. cited by applicant

Knoblauch et al., “Memory capacities for synaptic and structural plasticity,” Neural Comput., Feb. 2010, 22:289-341. cited by applicant

Kulakov, “Features of a Simple Psychophysiological Reaction,” Human Physiology, Jun. 15, 2017, 44:412-417. cited by applicant

Latora et al., “Efficient behavior of small-world networks,” Phys. Rev. Lett., Oct. 17, 2001, 87:1-4. cited by applicant

Le Be et al., “Spontaneous and evoked synaptic rewiring in the neonatal neocortex,” Proc. Natl. Acad. Sci. U.S.A., Aug. 29, 2006 103:13214-13219. cited by applicant

Levi, “A short course on Algebraic topology geared towards applications to Neuroscience,” University of Aberdeen, Jul. 18, 2017, pp. 1-14. cited by applicant

Luczak et al., “Packetbased communication in the cortex,” Nat. Rev. Neurosci., Oct. 28, 2015, 16:745-755. cited by applicant

Markram et al., “Reconstruction and simulation of neocortical microcircuitry,” Cell, Oct. 8, 2015, 163:456-492. cited by applicant

McCoss, Agency of Life, Entropic Gravity and Phenomena Attributed to “Dark Matter,” Journal of Quantum Information Science, Jun. 15, 2017, 7:67-75. cited by applicant

McCoss, “Lithium Quantum Consciousness,” Journal of Quantum Information Science, Nov. 8, 2017, 7:125-139. cited by applicant

Meyer et al., “Cell type-specific thalamic innervation in a col. of rat vibrissal cortex,” Cereb. Cortex, Jun. 9, 2010, 20: 2287-2303. cited by applicant

Miller et al., “Visual stimuli recruit intrinsically generated cortical ensembles,” Proc. Natl. Acad. Sci. U.S.A., Sep. 8, 2014, pp. E4053-E4061. cited by applicant

Miura et al., "Sparse Parallel Algorithms for Recognizing Touch Topology on Curved Interactive Screens," IEEE Access, Jul. 31, 2017, 5:14889-14897. cited by applicant

Munkres, "Elements of Algebraic Topology," Massachusetts Institute of Technology, 1984, Addison-Wesley Publishing Company, 233 pages. cited by applicant

Opris et al., "What Is the Evidence for Inter-laminar Integration in a Prefrontal Cortical Minicolumn?," Frontiers in Neuroanatomy, Dec. 14, 2017, 11: 1-11. cited by applicant

Pajevic et al., "The organization of strong links in complex networks," Nat. Phys., Mar. 11, 2012, 8:429-436. cited by applicant

Palm et al., "On the significance of correlations among neuronal spike trains," Biol. Cybern., Jun. 1, 1988, 59:1-11. cited by applicant

Patel et al., "Generative Design of Electromagnetic Structures Through Bayesian Learningm," IEEE Transactions on Magnetics, Oct. 20, 2017, 54:1-4. cited by applicant

Pedregosa et al., "Scikit-learn: machine learning in Python," J. Mach. Learn. Res., Oct. 2010, 12:2825-2830. cited by applicant

Perin et al., "A synaptic organizing principle for cortical neuronal groups," Proc. Natl. Acad. Sci. U.S.A., Mar. 29, 2011, 108:5419-5424. cited by applicant

Peters et al., "The projection of the lateral geniculate nucleus to area 17 of the rat cerebral cortex. I General description.," J. Neurocytol., Feb. 1976, 5:63-84. cited by applicant

Petlevski, "Modeling the Model: the World Beyond the Immediate Sensorium," Proceedings of the 20th International Multiconference Information Society, Oct. 18, 2017, pp. 44-47. cited by applicant

Petri et al., "Homological scaffolds of brain functional networks," J. R. Soc. Interface, Dec. 6, 2014, 11:1-11. cited by applicant

Planas et al., "Accelerating Data Analysis in Simulation Neuroscience with Big Data Technologies," International Conference on Computational Science, Jun. 1, 2018, pp. 363-377. cited by applicant

Ramaswamy et al., "The neocortical microcircuit collaboration portal: a resource for rat somatosensory cortex," Front. Neural Circuits, Oct. 8, 2015, 9:1-14. cited by applicant

Reimann et al. "An algorithm to predict the connectome of neural microcircuits," Front. Comput. Neurosci., Oct. 8, 2015, 9:1-18. cited by applicant

Renart et al., "The asynchronous state in cortical circuits," Science 327, Jan. 29, 2010, 587-590. cited by applicant

Rosenbaum et al., "The spatial structure of correlated neuronal variability," Nat. Neurosci., Oct. 21, 2016, 20:107-114. cited by applicant

SciPy.org [online], " SciPy: Open Source Scientific Tools for Python," available on or before Mar. 9, 2001, via Internet Archive: Wayback Machine URL
<https://web.archive.org/web/20010309040805/http://www.scipy.org/>, retrieved on Jan. 17, 2020,
<https://www.scipy.org/citing.html>. cited by applicant

See et al., "Coordinated neuronal ensembles in primary auditory cortical columns," Elife Sciences, Jun. 5, 2018, pp. 1-33. cited by applicant

Shepherd et al., "Geometric and functional organization of cortical circuits," Nat. Neurosci., May 8, 2005, 8:782-790. cited by applicant

Song et al. "Highly nonrandom features Of synaptic connectivity in local cortical circuits," PLoS Biol., Mar. 1, 2005, 3:0507-0519. cited by applicant

Stepanyants et al., "Neurogeometry and potential synaptic connectivity," Trends in Neurosci., Jun. 2, 2005, 28:387-394. cited by applicant

Suarez, "The interaction between task goals and the representation of choice options in decision-making," Thesis for the degree of Doctor of Philosophy, University College of London, Sep. 2017, pp. 1-176. cited by applicant

Tozzi et al., "Brain Projective Reality: Novel Clothes for the Emperor, Reply to comments of "Topodynamics of metastable brains," by Tozzi et al." Physics of Life Reviews, Jun. 28, 2017, pp. 1-11. cited by applicant

Tyukin et al., "High-Dimensional Brain: A Tool for Encoding and Rapid Learning of Memories by Single Neurons," *Bull Math Biol.*, Mar. 19, 2018, 11:4856-4888. cited by applicant

Varshney et al., "Structural properties of the caenorhabditis elegans neuronal network," *PLoS Comput. Biol.*, Feb. 3, 2011, 7:1-22. cited by applicant

Wijaya et al., "Finding an appropriate equation to measure similarity between binary vectors: case studies on Indonesian and Japanese herbal medicines," *BMC bioinformatics*, Dec. 2016, 17:1-19. cited by applicant

Willshaw et al., "Nonholographic associative memory," *Nature* 222, Jun. 7, 1969, pp. 960-963. cited by applicant

Guerreiro et al., "A Neural Key Generator for a Public Block Cipher," *IEEE Ninth Brazilian Symposium on Neural Networks (SBRN '06)*, Oct. 23-27, 2006, Ribeirao Preto, BR, 143-147. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/056588, dated May 26, 2020, 12 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/056686, dated May 26, 2020, 13 pages. cited by applicant

Bahraini et al., "Topological pattern selection in recurrent networks," *Neural Networks*, Feb. 9, 2012, 31:22-32. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085007, dated Mar. 24, 2021, 15 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085016, dated Mar. 24, 2021, 18 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085169, dated Mar. 18, 2021, 15 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085200, dated Apr. 4, 2021, 15 pages. cited by applicant

PCT International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085434, dated Mar. 24, 2021, 16 pages. cited by applicant

Salnikov et al., "Simplicial complexes and complex systems," *arxiv.org*, Cornell University Library, Jul. 20, 2018, 8 pages. cited by applicant

Yang et al., "Road Detection and Centerline Extraction via Deep Recurrent Convolutional Neural Network U-Net," *IEEE Transactions on Geoscience and Remote Sensing*, Sep. 2019, 57(9):7209-7220. cited by applicant

Cai et al., "Path-level network transformation for efficient architecture search," *In International Conference on Machine Learning*, Jul. 3, 2018 PMLR, 10 pages. cited by applicant

Ludermir et al., "An optimization methodology for neural network weights and architectures," *IEEE Transactions on Neural Networks*, Nov. 13, 2006, 17(6):1452-9. cited by applicant

Rawal et al., "From nodes to networks: Evolving recurrent neural networks," *arXiv preprint arXiv:1803.04439*, Mar. 12, 2018, 8 pages. cited by applicant

Lena et al., "Deep spatio-temporal architectures and learning for protein structure prediction," *Advances in neural information processing systems*, 25, 2012, 9 pages. cited by applicant

Park et al., "CNN (Convolutional Neural Network) based in-loop filter in HEVC," *In Proceedings of the Korean Society of Broadcast Engineers Conference 2016*, The Korean Institute of Broadcast and Media Engineers, 2016, pp. 369-372 (with English Abstract). cited by applicant

Wang et al., "Topological recurrent neural network for diffusion prediction," *In 2017 IEEE International Conference on Data Mining (ICDM)*, Nov. 18, 2017, pp. 475-484. cited by applicant

Masulli et al., "Dynamics of evolving feed-forward neural networks and their topological invariants," *Network and Parallel Computing*, Aug. 13, 2016, pp. 99-106. cited by applicant

Bauer et al., "Real-time ultra-low power ECG anomaly detection using an event-driven neuromorphic processor," *IEEE transactions on biomedical circuits and systems*, Nov. 11, 2019, 13(6):1575-82. cited by applicant

Goodman et al., "Brian: a simulator for spiking neural networks in python," *Frontiers in neuroinformatics*, 2008:5, 10 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085007, dated Jun. 23, 2022, 8 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085016, dated Jun. 23, 2022, 11 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085169, dated Jun. 23, 2022, 8 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085200, dated Jun. 23, 2022, 9 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085434, dated Jun. 23, 2022, 9 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085716, dated Jun. 23, 2022, 11 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085716, dated Mar. 25, 2021, 18 pages. cited by applicant

Minkovich et al. "Programming time-multiplexed reconfigurable hardware using a scalable neuromorphic compiler," *IEEE transactions on neural networks and learning systems*, Apr. 11, 2012, 23(6):889-901. cited by applicant

U.S. Appl. No. 17/783,961, filed Jun. 9, 2022, Markram. cited by applicant

Andersen et al., "Overlapping clusters for distributed computation," In *Proceedings of the Fifth ACM International Conference on Web Search and Data Mining*, Feb. 8, 2012, 9 pages. cited by applicant

Baskararaja et al., "Subgraph matching using graph neural network," *Journal of Intelligent Learning Systems and Applications*, Nov. 28, 2012, 4(04):274-278. cited by applicant

Bauer et al., "Clear and Compress: Computing Persistent Homology in Chunks," *Topological Methods in Data Analysis and Visualization III.*, 2014, pp. 103-117. cited by applicant

Cormen et al., "Introduction to Algorithms," Copyright 2001 by The Massachusetts Institute of Technology, 984 pages. cited by applicant

Erdos and Renyi, "On random graphs, I," *Math. Debrecen.*, 1959, 6:290-297. cited by applicant

Gros, "Cognitive computation with autonomously active neural networks: an emerging field," *Cognitive Computation*, Mar. 2009, 1(1):77-90. cited by applicant

Gros, "Neural networks with transient state dynamics," *New Journal of Physics*, Apr. 30, 2007, 9(4):109, 21 pages. cited by applicant

Gros, "Self-sustained thought processes in a dense associative network," In *Annual Conference on Artificial Intelligence*, Sep. 11, 2005, Springer, Berlin, Heidelberg, 14 pages. cited by applicant

Hatcher, "Algebraic Topology," Cambridge University Press, Feb. 2002, 556 pages. cited by applicant

Kahle, "Sharp vanishing thresholds for cohomology of random flag complexes," *Ann. of Math.*, May 2014, 179(3):1085-1107. cited by applicant

Lin et al., "Organizing principles of real-time memory encoding: neural clique assemblies and universal neural codes," *TRENDS in Neurosciences*, Jan. 2006, 29(1):48-57. cited by applicant

Maria, "Algorithms and data structures in computational topology" (Doctoral dissertation, Université Nice Sophia Antipolis), Oct. 28, 2014, 205 pages. cited by applicant

mathworld.wolfram.com [online], "Adjacency Matrix," 2016, retrieved via Internet Archive on Apr. 8, 2022, retrieved from <<https://web.archive.org/web/20160311204902/https://mathworld.wolfram.com/AdjacencyMatrix.htm1>>, 2 pages. cited by applicant

Milo et al., "Network Motifs: Simple Building Blocks of Complex Networks," *Science*, Oct. 25, 2002, 298(5594):824-827. cited by applicant

Peters et al., "The projection of the lateral geniculate nucleus to area 17 of the rat cerebral cortex. V. Degenerating axon terminals synapsing with Golgi impregnated neurons," *J Neurocytol.*, 1979, 8:331-

357. cited by applicant

Srivastava et al., "Dropout: a simple way to prevent neural networks from overfitting," The journal of machine learning research, Jan. 2014, 15(1):1929-58. cited by applicant

TW Office Action in Taiwanese Appln. No. 10943863, dated Mar. 1, 2022, 13 pages (with machine translation). cited by applicant

U.S. Appl. No. 16/004,635, filed Jun. 11, 2018, Henry Markram. cited by applicant

U.S. Appl. No. 16/004,837, filed Jun. 11, 2018, Henry Markram. cited by applicant

U.S. Appl. No. 16/004,757, filed Jun. 11, 2018, Henry Markram. cited by applicant

U.S. Appl. No. 16/004,671, filed Jun. 11, 2018, Henry Markram. cited by applicant

U.S. Appl. No. 16/356,391, filed Mar. 18, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 16/356,478, filed Mar. 18, 2019, Kathryn Hess. cited by applicant

U.S. Appl. No. 16/710,266, filed Dec. 11, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 16/710,058, filed Dec. 11, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 16/710,205, filed Dec. 11, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 16/710,176, filed Dec. 11, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 16/710,080, filed Dec. 11, 2019, Henry Markram. cited by applicant

U.S. Appl. No. 15/864,146, filed Jan. 8, 2018, Michael Wolfgang Reimann. cited by applicant

Baker et al., "Designing neural network architectures using reinforcement learning," arXiv preprint, arXiv:1611.02167, Nov. 7, 2016, 18 pages. cited by applicant

Chambers et al., "Higher-order synaptic interactions coordinate dynamics in recurrent networks," PLoS computational biology, Aug. 19, 2016, 12(8):e1005078, 23 pages. cited by applicant

EP Extended European Search Report in EP Appln. No. 17174314.9, dated Dec. 14, 2017, 11 pages. cited by applicant

Garcia et al., "Unsupervised classification of neural spikes with a hybrid multilayer artificial neural network," Journal of Neuroscience Methods, 1998, 82:59-73. cited by applicant

Giusti et al., "Two's company, three (or more) is a simplex," Journal of computational neuroscience, Aug. 2016, 41(1):1-14. cited by applicant

Graves, "Adaptive computation time for recurrent neural networks," arXiv preprint arXiv: 1603.08983, Mar. 29, 2016, 19 pages. cited by applicant

Gros et al., "Semantic learning in autonomously active recurrent neural networks," Logic Journal of IGPL, Oct. 1, 2010, 18(5):686-704. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/056588, dated Sep. 30, 2021, 8 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/056686, dated Sep. 30, 2021, 10 pages. cited by applicant

Sizemore et al., "Cliques and cavities in the human connectome," Journal of computational neuroscience, Feb. 2018, 44(1):115-45. cited by applicant

Achard et al., "A Resilient, Low-Frequency, Small-World Human Brain Functional Network with Highly Connected Association Cortical Hubs," Journal of Neuroscience, Jan. 4, 2006, 26(1):63-72. cited by applicant

Arai et al., "The Effects of Theta Precession on Spatial Learning and Simplicial Complex Dynamics in a Topological Model of the Hippocampal Spatial Map," PLOS Computational Biology, Jun. 19, 2014, 10(6):r1003651. cited by applicant

Astrivis [online], "Demo Scan Church Dec. 2016," Dec. 2, 2016, retrieved on Nov. 27, 2022, retrieved from URL <<https://www.youtube.com/watch?v=xCGQ4aaB1Go>>, 2 pages [Video Submission]. cited by applicant

Bassett et al., "Dynamic reconfiguration of human brain networks during learning," PNAS, Apr. 18, 2011, 108(18):7641-7646. cited by applicant

Brown et al., "Structure of the Afferent Terminals in Terminal Ganglion of a Cricket and Persistent Homology," PLOS One, May 23, 2012, 7(5):e37278. cited by applicant

Chen et al., "Neural Representation of Spatial Topology in the Rodent Hippocampus," *Neural Computation*, Jan. 2014, 26(1):1-39. cited by applicant

Choi et al., "Abnormal metabolic connectivity in the pilocarpine-induced epilepsy rat model: A multiscale network analysis based on persistent homology," *NeuroImage*, Oct. 1, 2014, 99:226-236. cited by applicant

Chung et al., "Persistence Diagrams of Cortical Surface Data," *Information Processing in Medical Imaging*, 2009, pp. 386-397. cited by applicant

Crossley et al., "Cognitive relevance of the community structure of the human brain functional coactivation network," *PNAS*, Jun. 24, 2013, 110(28):11583-11588. cited by applicant

Dongjiang et al., "DXSLAM: A Robust and Efficient Visual SLAM System with Deep Features," *ArXiv*, Aug. 12, 2020, 8 pages. cited by applicant

Drakesmith et al., "Overcoming the effects of false positives and threshold bias in graph theoretical analyses of neuroimaging data," *NeuroImage*, Sep. 2015, 118:313-333. cited by applicant

Ellis et al., "Describing High-Order Statistical Dependence Using "Concurrence Topology," With Application to Functional MRI Brain Data," *Homology, Homotopy and Applications*, 2014, 16(1):245-264. cited by applicant

Garrison et al., "The (in)stability of functional brain network measures across thresholds," *NeuroImage*, Sep. 2015, 118:651-661. cited by applicant

Ginestet et al., "Brain Network Analysis: Separating Cost from Topology Using Cost-Integration," *PLOS One*, Jul. 28, 2018, 2018, 13(7):e21570. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2022/073852, mailed on Dec. 9, 2022, 18 pages. cited by applicant

Khalid et al., "Tracing the evolution of multi-scale functional networks in a mouse model of depression using persistent brain network homology," *NeuroImage*, Nov. 1, 2014, 101:351-363. cited by applicant

Kim et al., "Morphological brain network assessed using graph theory and network filtration in deaf adults," *Hearing Research*, Sep. 2014, 315:88-98. cited by applicant

Langer et al., "The Problem of Thresholding in Small-World Network Analysis," *PLOS One*, Jan. 3, 2013, 8(1):e53199. cited by applicant

Lee et al., "Discriminative persistent homology of brain networks," *2011 IEEE International Symposium on Biomedical Imaging: From Nano to Macro*, Mar. 30, 2011, pp. 841-844. cited by applicant

Lucky Robot [online], "ORB SLAM3: Ar demo," Feb. 23, 2021, retrieved on Nov. 27, 2022, retrieved from URL <https://www.youtube.com/watch?v=c1ExX_IA5tM>, 2 pages [Video Submission]. cited by applicant

Medaglia et al., "Cognitive Network Neuroscience," *Journal of Cognitive Neuroscience*, Aug. 1, 2015, 27(8):1471-1491. cited by applicant

Mur-Artal et al., "ORB-SLAM2: an Open-Source SLAM System for Monocular, Stereo and RGB-D Cameras," *IEEE Transactions on Robotics*, Oct. 5, 2015, 31(5):1147-1163. cited by applicant

Pirino et al., "A topological study of repetitive co-activation networks in in vitro cortical assemblies," *Physical Biology*, Jan. 5, 2015, 12(1):016007. cited by applicant

Singh et al., "Topological analysis of population activity in visual cortex," *Journal of Vision*, Jun. 2008, 8(8):11, 18 pages. cited by applicant

Stolz, "Computational Topology in Neuroscience," *Dissertation for the Degree of Master of Science in Mathematical Modelling & Scientific Computing at the University of Oxford*, Sep. 2014, 77 pages. cited by applicant

Alaniz, "Machiavellian Playbook for Artificial General Intelligence (AGI), Early Doctrine for AGI power," *Research Gate*, Aug. 2018, pp. 1-42. cited by applicant

Alaniz, "Remaking Einstein with a Dynamical Graph Rich Person Object Avatar Extraction Markup Language (intelligence as topological algebraic invariants, graph entropies, and dynamics)," Mar. 16, 2019, 14 pages. cited by applicant

Allswede et al., "Prenatal inflammation and risk for schizophrenia: A role for immune proteins in

neurodevelopment,” *Development and Psychopathology*, Aug. 2, 2018, 30: 1157-1178. cited by applicant

Antonopoulos et al., “Evaluating performance of neural codes in neural communication networks,” *Neural Networks*, Sep. 24, 2018, pp. 1-17. cited by applicant

Baptiste-Bardin et al., “Topological exploration of artificial neuronal network dynamics,” *Network Neuroscience*, Jan. 24, 2019, pp. 1-28. cited by applicant

Basset et al., “Network models in neuroscience,” *arXiv*, Jul. 21, 2018, pp. 1-12. cited by applicant

Basset et al., “On the nature and use of models in network neuroscience,” *Nature Reviews*, Jul. 12, 2018, 19:566-578. cited by applicant

Baudot et al., “Topological Information Data Analysis,” *Entropy*, Sep. 2019, 21:1-38. cited by applicant

Betzel, “Stability of spontaneous, correlated activity in mouse auditory cortex,” *PLOS: Computational Biology*, Dec. 9, 2019, 1-25. cited by applicant

Bianconi et al., “Topological percolation on hyperbolic simplicial complexes,” *Phys. Rev. E*, Nov. 21, 2018, 5:1-18. cited by applicant

Busch et al., “Intrinsic temporal tuning of neurons in the optic tectum is shaped by multisensory experience,” *Journal of Neurophysiology*, Sep. 5, 2019, pp. 1-23. cited by applicant

Chacholski et al., “On the topology of complexes of injective words,” *Journal of Applied and Computational Topology*, Aug. 16, 2019, 21:1-16. cited by applicant

Chowdhury et al., “Path homologies of deep feedforward networks,” *Arxiv*, Oct. 16, 2019, pp. 1-6. cited by applicant

Coggan et al., “A Process for Digitizing and Simulating Biologically Realistic Oligocellular Networks Demonstrated for the Neuro-Glio-Vascular Ensemble,” *Neurosci*, Sep. 25, 2018, 12:1-21. cited by applicant

Collins, “The case for emulating insect brains using anatomical “wiring diagrams” equipped with biophysical models of neuronal activity,” *Biological cybernetics*, Dec. 1, 2019, 113: 465-75. cited by applicant

Curto et al., “Cell groups reveal structure of stimulus space,” *Plos Comput. Biol.*, Oct. 2008, 4(10):e100205. cited by applicant

Curto, “Relating network connectivity to dynamics: opportunities and challenges for theoretical neuroscience,” *Current Opinion in Neurobiology*, Oct. 1, 2019, 58:11-20. cited by applicant

Curto, “What Can Topology Tell US About the Neural Code?,” *The American Mathematical Society*, Jan. 2017, 54(1):63-78. cited by applicant

Dabaghian et al., “A topological paradigm for hippocampal spatial map formation using persistent homology,” *Plos Comput. Biol.*, Aug. 2012, 8(8):e1002581. cited by applicant

Dabaghian et al., “Reconceiving the hippocampal map as a topological template,” *Elife*, Aug. 20, 2014, 3:e03476. cited by applicant

DeCharms et al., “Primary cortical representation of sounds by the coordination of action-potential timing,” *Nature*, Jun. 13, 1996, 381(6583):610-613. cited by applicant

Doborjeh et al., “Spiking Neural Network Modelling Approach Reveals How Mindfulness Training Rewires the Brain,” *Scientific Reports*, Apr. 23, 2019, 9:1-15. cited by applicant

Ebli, “A Notion of Harmonic Clustering in Simplicial Complexes,” *Laboratory for Topology and Neuroscience, École Polytechnique Fédérale de Lausanne*, 1015 Lausanne, Switzerland, Oct. 17, 2019, 9 pages. cited by applicant

Fan et al., “A Brief History of Simulation Neuroscience,” *Front. Neuroinform.*, May 7, 2019, 13:1-28. cited by applicant

Feldbrugge et al., “Stochastic homology of Gaussian vs. non-Gaussian random fields: graphs towards Betti Nos. and persistence diagrams,” *Journal of Cosmology and Astroparticle Physics*, Sep. 24, 2019, 49 pages. cited by applicant

Feng et al., “Persistent Homology of Geospatial Data: A Case Study with Voting,” Jan. 30, 2019, pp. 1-29. cited by applicant

Garcia et al., “Reconfigurations within resonating communities of brain regions following TMS reveal

different scales of processing, ” BioRxiv, Jan. 9, 2020, pp. 1-39. cited by applicant

Guerreno, “The sleep onset transition: a connectivity investigation built on EEG source localization,” University of Zurich, Facult of Science, Aug. 6, 2018, pp. 1-292. cited by applicant

Hu et al., “Local paths to global coherence: cutting networks down to size,” Phys. Rev. E, Mar. 10, 2014, 89: 1-16. cited by applicant

Ibanze-Marcelo et al., “Topology highlights mesoscopic functional equivalence between imagery and perception: The case of hypnotizability,” NeuroImage, Jun. 19, 2019, 437-449. cited by applicant

Ignacio et al., “Classification of Single-lead Electrocardiograms: TDA Informed Machine Learning,” ArXiv, Nov. 25, 2019, pp. 1-6. cited by applicant

Ivancevic et al., “Tensor-Centric Warfare V: Topology of Systems Confrontation,” Intelligent Control Automation, Feb. 28, 2019, 10:13-45. cited by applicant

Jozefowicz et al., “An Empirical Exploration of Recurrent Network Architectures,” Proceedings of the 32nd International Conference on Machine Learning, 2015, 37:1-9. cited by applicant

Kanari et al., “A Topological Representation of Branching Neuronal Morphologies,” Neuroinformatics, Oct. 3, 2017, 11 pages. cited by applicant

Kanari et al., “Quantifying topological invariants of neuronal morphologies,” ArXiv, Mar. 28, 2016, 15 pages. cited by applicant

Kartun-Giles, “Beyond the clustering coefficient: A topological analysis of node neighborhoods in complex networks,” Chaos, Solitons & Fractals: X, Feb. 16, 2019, 1:1-12. cited by applicant

Kastanenka et al., “A roadmap to integrate astrocytes into Systems Neuroscience,” Gila, Wiley Periodicals, Apr. 9, 2019, pp. 1-22. cited by applicant

Khajezade et al., “A Game-Theoretical Network Formation Model for C. elegans Neural Network,” Frontiers in Computational Neuroscience, Jul. 9, 2019, 13:1-12. cited by applicant

Kumarashinghe, “Deep learning and deep knowledge representation in Spiking Neural Networks for Brain-Computer Interface,” Neural Networks, Sep. 20, 2019, 121:169-185. cited by applicant

Kumbhar et al., “CoreNeuron: An Optimized Compute Engine for the NEURON Simulator,” Frontiers in Neuroinformatics, Sep. 4, 2019, pp. 1-27. cited by applicant

Kvam, “A geometric framework for modeling dynamic decisions among arbitrarily many alternatives,” Journal of Mathematical Psychology, Aug. 1, 2019, 91:14-37. cited by applicant

Luccioli et al., “Modeling driver cells in developing neural networks,” PLOS Computational Biology, Nov. 2, 2018, pp. 1-31. cited by applicant

Lutgehetmann et al., “Computing Persistent Homology of Directed Flag Complexes,” Algorithms, Jan. 7, 2020, 1:1-18. cited by applicant

Lynn et al., “The physics of brain network structure, function, and control,” Nature Reviews, May 27, 2019, 1:318-332. cited by applicant

Mardones, “Persistence Steenrod modules,” ArXiv Mathematics, Dec. 12, 2018, pp. 1-10. cited by applicant

Mehta, “Storing and retrieving long-term memories: cooperation and competition in synaptic dynamics,” Advances in Physics: X, Jul. 19, 2018, 3:756-790. cited by applicant

Millán et al., “Complex Network Geometry and Frustrated Synchronization,” Scientific Reports, Jul. 2, 2018, 8:1-10. cited by applicant

Millan et al., “Synchronization in network geometries with finite spectral dimension,” ArXiv, Jan. 31, 2019, pp. 1-15. cited by applicant

Muller et al., “Neocortical plasticity: an unsupervised cake but No. free lunch,” 33rd Conference on Neural Information Processing Systems, Vancouver, Canada, Nov. 2019, 6 pages. cited by applicant

Nolte, “Cortical reliability amid noise and chaos,” Nature Communications, Aug. 22, 2019, 10:1-15. cited by applicant

Norton, “Dynamics of Reaction-Diffusion Oscillators in Star and other Networks with Cyclic Symmetries Exhibiting Multiple Clusters,” Physical Review Letters, Oct. 4, 2019, 123: 1-6. cited by applicant

Office Action in U.S. Appl. No. 16/004,671, dated Oct. 27, 2022, 40 pages. cited by applicant

Office Action in US Appln. No. 16/710,2056, dated Oct. 26, 2022, 51 pages. cited by applicant

Peng, “High-throughput microcircuit analysis of individual human brains through next-generation multineuron patch-clamp,” *BioRxiv*, Jan. 1, 2019, pp. 1-49. cited by applicant

Petri et al., “Simplicial Activity Driven Model,” *Phys. Rev. Lett.*, Nov. 29, 2018, 121:1-5. cited by applicant

Popa, “Psychology 2.0: The Emergence of Individuality,” Sep. 2019, pp. 1-6. cited by applicant

Prentner, “Consciousness and Topologically Structured Phenomenal Spaces,” *Consciousness and Cognition*, Feb. 26, 2019, 70:25-38. cited by applicant

Santos et al., “Topological phase transitions in functional brain networks,” *Physical Review E*, Sep. 30, 2019, 100: 1-17. cited by applicant

Saucan et al., “Forman's Ricci Curvature—From Networks to Hypernetworks,” vol. 1, *Proceedings The 7th International Conference on Complex Networks and Their Applications*, ArXiv, Oct. 17, 2018, 13 pages. cited by applicant

Schoenberg et al., “Mapping Meditative States and Stages with Electrophysiology: Concepts, Classifications, and Methods,” *Current Opinion in Psychology*, Oct. 18, 2018, 28:211-217. cited by applicant

Sizemore et al., “The importance of the whole: Topological data analysis for the network neuroscientist,” *Network Neuroscience*, Oct. 17, 2018, 3:1-18. cited by applicant

Skardal et al., “Abrupt Desynchronization and Extensive Multistability in Globally Coupled Oscillator Simplexes,” *Physical Review Letters* 122, Jun. 19, 2019, pp. 1-6. cited by applicant

Timsit et al., “Nervous-Like Circuits in the Ribosome Facts, Hypotheses and Perspectives,” *Int. J. Mol. Sci*, Jun. 14, 2019, 20:1-22. cited by applicant

Tran et al., “Scale-variant topological information for characterizing the structure of complex networks,” *Physical Review E*. Sep. 18, 2019, 100:1-18. cited by applicant

Turner, “Rips filtrations for quasimetric spaces and asymmetric functions with stability results,” *Algebraic & Geometric Topology*, May 21, 2019, 19:1135-1170. cited by applicant

Velazquez et al., “On a Simple General Principle of Brain Organization,” *Frontiers in Neuroscience*, Oct. 15, 2019, 13:1-16. cited by applicant

Young, “Standard Laterality Models: Neo-Eriksonian Perspectives,” Chapter 8, *Causality and Development*, Jan. 2019, pp. 147-179. cited by applicant

Office Action in U.S. Appl. No. 16/710,266, dated Nov. 17, 2022, 32 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,635, dated Jan. 23, 2023, 24 pages. cited by applicant

Djurfeldt et al., “Efficient generation of connectivity in neuronal networks from simulator-independent descriptions,” *Fronteriers in Neuroinformatics*, Mar. 28, 2014, 8(43):1-11. cited by applicant

docs.opencv.org [online], “Camera Calibration and 3D Reconstruction,” Apr. 14, 2020, retrieved on Oct. 5, 2022, retrieved from URL <https://docs.opencv.org/4.4.0/d9/d0c/group_calib3d.html#ga549c2075fac14829ff4a58bc931c033d>, 78 pages. cited by applicant

docs.opencv.org [online], “Pose Estimation,” available no later than Sep. 13, 2021, retrieved on Oct. 5, 2022, retrieved from URL <https://docs.opencv.org/4.x/d7/d53/tutorial_py_pose.html>, 3 pages. cited by applicant

Extended European Search Report in European Appln. No. 17174316.4, dated Dec. 14, 2017, 11 pages. cited by applicant

Extended European Search Report in European Appln. No. 17174317.2, dated Dec. 14, 2017, 11 pages. cited by applicant

Github.com [online], “facebookresearch/detectron2,” 2019, retrieved on Oct. 5, 2022, retrieved from URL <<https://github.com/facebookresearch/detectron2>>, 4 pages. cited by applicant

Gleeson et al., “Open Source Brain: A Collaborative Resource for Visualizing, Analyzing, Simulating, and Developing Standardized Models of Neurons and Circuits,” *Neuron*, Aug. 7, 2019, 103(3):395-

411.e5. cited by applicant

Gripon et al., "Sparse Neural Networks with Large Learning Diversity," IEEE Transactions on Neural Networks, Jul. 2011, 22(7):1087-1096. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085750, mailed on Jun. 23, 2022, 8 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085754, mailed Jun. 23, 2022, 12 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2020/085762, mailed on Jun. 23, 2022, 7 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085716, mailed on Mar. 25, 2021, 13 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085750, mailed on Apr. 6, 2021, 10 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085754, mailed Apr. 6, 2021, 14 pages. cited by applicant

International Search Report and Written Opinion in International Appln. No. PCT/EP2020/085762, mailed on Apr. 6, 2021, 9 pages. cited by applicant

Liang et al., "Interpretable Structure-Evolving LSTM," 2017 IEEE Conference on Computervision and Pattern Recognition (CVPR), 2017, pp. 2175-2184. cited by applicant

Office Action in Korean Appln. No. 2020/7035845, dated Feb. 16, 2022, 9 pages (with English translation). cited by applicant

Office Action in Korean Appln. No. 2020-7035843, dated Jan. 27, 2022, 9 pages (with English translation). cited by applicant

Office Action in Korean Appln. No. 2020-7035844, dated Feb. 17, 2022, 6 pages (with English translation). cited by applicant

Office Action in Korean Appln. No. 2020-7035846, dated Feb. 15, 2022, 6 pages (with English translation). cited by applicant

Office Action in Korean Appln. No. 2020-7035847, dated Aug. 10, 2022, 5 pages (with English translation). cited by applicant

Office Action in Korean Appln. No. 2020-7035847, dated Feb. 15, 2022, 11 pages (with English translation). cited by applicant

Office Action in Taiwanese Appln. No. 108119813, dated May 19, 2021, 8 pages (with machine translation). cited by applicant

Office Action in Taiwanese Appln. No. 109143238, dated Dec. 1, 2021, 18 pages (with machine translation). cited by applicant

Office Action in Taiwanese Appln. No. 109143239, dated Nov. 30, 2021, 22 pages (with machine translation). cited by applicant

Office Action in Taiwanese Appln. No. 109143240, dated Dec. 24, 2021, 13 pages (with machine translation). cited by applicant

Office Action in Taiwanese Appln. No. 109143242, dated Nov. 26, 2021, 24 pages (with machine translation). cited by applicant

Office Action in Taiwanese Appln. No. 109143863, dated Mar. 1, 2022, 11 pages (with machine translation). cited by applicant

Office Action in U.S. Appl. No. 15/864,146, dated Dec. 8, 2021, 73 pages. cited by applicant

Office Action in U.S. Appl. No. 15/864,146, dated Jan. 13, 2021, 37 pages. cited by applicant

Office Action in U.S. Appl. No. 15/864,146, dated Jun. 3, 2021, 55 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,635, dated Aug. 3, 2022, 24 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,635, dated Jun. 3, 2021, 44 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,635, dated Nov. 24, 2021, 34 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,671, dated May 26, 2022, 43 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,671, dated Nov. 17, 2021, 40 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,757, dated Dec. 20, 2021, 31 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,757, dated Jul. 14, 2022, 30 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,757, dated Jun. 8, 2021, 26 pages. cited by applicant

Office Action in U.S. Appl. No. 16/004,837, dated Jun. 6, 2022, 8 pages. cited by applicant

Office Action in U.S. Appl. No. 16/356,391, dated Mar. 8, 2022, 10 pages. cited by applicant

Office Action in U.S. Appl. No. 16/356,391, dated Nov. 9, 2021, 13 pages. cited by applicant

Office Action in U.S. Appl. No. 16/356,478, dated Mar. 3, 2022, 5 pages. cited by applicant

Office Action in U.S. Appl. No. 16/710,058, dated Sep. 9, 2022, 14 pages. cited by applicant

Office Action in U.S. Appl. No. 16/710,176, dated Jun. 9, 2022, 7 pages. cited by applicant

Panda et al., "Learning to Recognize Actions from Limited Training Examples Using a Recurrent Spiking Neural Model", arXiv, Oct. 19, 2017, 17 pages. cited by applicant

Paredes-Valles et al., "Unsupervised Learning of a Hierarchical Spiking Neural Network for Optical Flow Estimation: From Events to Global Motion Perception," IEEE Transactions on Pattern Analysis and Machine Intelligence, Mar. 5, 2019, 42(8):2051-2064. cited by applicant

wikipedia.org [online], "DBSCAN," Mar. 31, 2008, retrieved on Oct. 5, 2022, retrieved from URL <<https://en.wikipedia.org/wiki/DBSCAN>>, 8 pages. cited by applicant

wikipedia.org [online], "Harris corner detector," Apr. 21, 2019, retrieved on Oct. 5, 2022, retrieved on Oct. 5, 2022, retrieved from URL <https://en.wikipedia.org/wiki/Harris_corner_detector>, 6 pages. cited by applicant

wikipedia.org [online], "OPTICS algorithm," Jun. 12, 2010, retrieved Oct. 5, 2022, retrieved from URL <https://en.wikipedia.org/wiki/OPTICS_algorithm>, 6 pages. cited by applicant

wikipedia.org [online], "Scale-invariant feature transform," Apr. 15, 2005, retrieved on Oct. 5, 2022, retrieved from URL <https://en.wikipedia.org/wiki/Scale-invariant_feature_transform>, 19 pages. cited by applicant

Yang et al., "Convolutional Neural Networks with Alternately Updated Clique," 2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2018, pp. 2413-2422. cited by applicant

Ghosh et al., "Design and Construction of a Brain-like computer: A New class of Frequency-Fractal Computing Using Wireless Communication in a Supramolecular Organic, Inorganic System," Information, 2014, 5, 28-100: doi: 10.339/info5010028. cited by applicant

Masulli et al., "Algebro-topological invariants in network theory, " Int. J. Complex Systems in Science, Sep. 2015, 5(1):13-17. cited by applicant

Office Action in U.S. Appl. No. 16/004,757, dated Mar. 29, 2023, 37 pages. cited by applicant

Office Action in U.S. Appl. No. 16/710,266, dated Apr. 12, 2023, 42 pages. cited by applicant

Office Action in U.S. Appl. No. 16/710,205, dated May 10, 2023, 37 pages. cited by applicant

U.S. Appl. No. 18/295,959, filed Apr. 5, 2023, Henry. cited by applicant

Office Action in European Appln. No. 19728992.9, mailed on Jan. 24, 2023, 8 pages. cited by applicant

Office Action in European Appln. No. 19728993.7, mailed Jan. 24, 2023, 8 pages. cited by applicant

Office Action in European Appln. No. 20829805.9, mailed Sep. 14, 2023, 10 pages. cited by applicant

U.S. Appl. No. 18/480,927, filed Oct. 20, 2023, Markram et al. cited by applicant

Office Action in U.S. Appl. No. 16/004,671, dated Jun. 28, 2023, 39 pages. cited by applicant

International Preliminary Report on Patentability in International Appln. No. PCT/EP2022/073852, mailed on Mar. 28, 2024, 12 pages. cited by applicant

Mur-Artal et al., "ORB-SLAM: a versatile and accurate monocular SLAM system," IEEE transactions on robotics, Aug. 24, 2015, 31(5):1147-63. cited by applicant

Mur-Artal et al., "ORB-SLAM2: An open-source slam system for monocular, stereo, and rgb-d cameras," IEEE transactions on robotics, Jun. 12, 2017, 33(5):1255-62. cited by applicant

Office Action in Chinese Appln. No. 202080035412.2, mailed on Jan. 30, 2024, 22 pages (with English translation). cited by applicant

Office Action in U.S. Appl. No. 16/710,266, dated Jan. 29, 2024, 35 pages. cited by applicant

Office Action in U.S. Appl. No. 18/161,414, dated Dec. 21, 2023, 33 pages. cited by applicant

Office Action in U.S. Appl. No. 18/167,958, dated Feb. 15, 2024, 7 pages. cited by applicant

U.S. Appl. No. 17/454,347, filed Nov. 10, 2021, Lütgebetmann. cited by applicant

Office Action in Chinese Appln. No. 201980053140.6, mailed on Apr. 29, 2024, 12 pages (with English translation). cited by applicant

Office Action in Chinese Appln. No. 201980053141.0, mailed on Apr. 30, 2024, 19 pages (with English translation). cited by applicant

Office Action in Chinese Appln. No. 201980053463.5, mailed on May 23, 2024, 7 pages (with English translation). cited by applicant

Office Action in Chinese Appln. No. 201980053465.4, mailed on May 23, 2024, 7 pages (with English translation). cited by applicant

Office Action in Chinese Appln. No. 201980054063.6, mailed on May 23, 2024, 7 pages (with English translation). cited by applicant

U.S. Appl. No. 16/004,635, filed Jun. 11, 2018, Markram. cited by applicant

U.S. Appl. No. 16/004,837, filed Jun. 11, 2018, Markram. cited by applicant

U.S. Appl. No. 16/004,757, filed Jun. 11, 2018, Markram. cited by applicant

U.S. Appl. No. 16/004,671, filed Jun. 11, 2018, Markram. cited by applicant

U.S. Appl. No. 18/611,781, filed Mar. 21, 2024, Markram. cited by applicant

U.S. Appl. No. 16/356,391, filed Mar. 18, 2019, Markram. cited by applicant

U.S. Appl. No. 18/295,959, filed Apr. 5, 2023, Markram. cited by applicant

U.S. Appl. No. 16/356,478, filed Mar. 18, 2019, Hess. cited by applicant

U.S. Appl. No. 18/161,414, filed Jan. 30, 2023, Hess. cited by applicant

U.S. Appl. No. 16/710,266, filed Dec. 11, 2019, Markram. cited by applicant

U.S. Appl. No. 16/710,058, filed Dec. 11, 2019, Markram. cited by applicant

U.S. Appl. No. 18/490,927, filed Oct. 20, 2023, Markram et al. cited by applicant

U.S. Appl. No. 16/710,205, filed Dec. 11, 2019, Markram. cited by applicant

U.S. Appl. No. 18/487,566, filed Oct. 16, 2023, Markram et al. cited by applicant

U.S. Appl. No. 16/710,176, filed Dec. 11, 2019, Markram. cited by applicant

U.S. Appl. No. 18/167,958, filed Feb. 13, 2023, Markram. cited by applicant

U.S. Appl. No. 16/710,080, filed Dec. 11, 2019, Markram. cited by applicant

U.S. Appl. No. 18/295,969, filed Apr. 5, 2023, Henry. cited by applicant

U.S. Appl. No. 18/656,669, filed May 7, 2024, Markram. cited by applicant

U.S. Appl. No. 17/783,976, filed Jun. 9, 2022, Markram. cited by applicant

U.S. Appl. No. 17/783,978, filed Jun. 9, 2022, Markram. cited by applicant

U.S. Appl. No. 17/783,981, filed Jun. 9, 2022, Markram. cited by applicant

U.S. Appl. No. 17/454,347, filed Nov. 10, 2021, Lütgehetmann. cited by applicant

U.S. Appl. No. 15/864,146, filed Jan. 8, 2018, Reimann. cited by applicant

U.S. Appl. No. 18/188,888, filed Mar. 23, 2023, Reimann. cited by applicant

Nazeran et al., "Neural Networks in Processing and Analysis of Biomedical Signals," Nonlinear Biomedical Signal Processing, Fuzzy Logic, Neural Networks, and New Algorithms, IEEE, Metin Akay (. ed), Jul. 2000, Chapter 4, pp. 69-97. cited by applicant

Office Action in U.S. Appl. No. 16/710,266, mailed on Aug. 9, 2024, 33 pages. cited by applicant

Timotheou, "The random neural network: a survey," The Computer Journal, Mar. 2010, 53(3):251-67. cited by applicant

U.S. Appl. No. 18/527,555, filed Dec. 4, 2023, Markram. cited by applicant

Google.com [online], "Search Query: Cliques simplices cavities," available on or before Feb. 19, 2025, retrieved on Feb. 19, 2025, retrieved from URL <<https://www.google.com/search?udm=36&q=cliques+simplices+cavities>>, 4 pages. cited by applicant

Office Action in Chinese Appln. No. 202080036778.1, mailed on Dec. 30, 2024, 26 pages (with English translation). cited by applicant

Office Action in U.S. Appl. No. 16/710,266, mailed on Feb. 13, 2025, 33 pages. cited by applicant
Office Action in U.S. Appl. No. 18/295,959, mailed on Nov. 18, 2024, 14 pages. cited by applicant
Office Action in U.S. Appl. No. 18/527,555, mailed on Nov. 21, 2024, 52 pages. cited by applicant
Chatzis et al., "The copula echo state network," Pattern Recognition, Jan. 2012, 45(1):570-7. cited by applicant

Khan et al., "Learning from Privacy Preserved Encrypted Data on Cloud Through Supervised and Unsupervised Machine Learning," Presented at Proceedings of the 2019 2nd International Conference on Computing, Mathematics and Engineering Technologies, Sukkur, Sindh, Pakistan, Jan. 30-31, 2019, 5 pages. cited by applicant

Office Action in Chinese Appln. No. 202080035412.2, mailed on Sep. 24, 2024, 9 pages (with Machine translation). cited by applicant

Office Action in European Appln. No. 20711136.0, mailed on Sep. 17, 2024, 8 pages. cited by applicant

Office Action in European Appln. No. 20711154.3, mailed on Sep. 13, 2024, 6 pages. cited by applicant

Office Action in U.S. Appl. No. 17/454,347, mailed on Oct. 16, 2024, 25 pages. cited by applicant

Office Action in Chinese Appln. No. 201980053465.4, mailed on Mar. 21, 2025, 11 pages (with English translation). cited by applicant

Office Action in U.S. Appl. No. 18/295,959, mailed on Mar. 5, 2025, 12 pages. cited by applicant

Primary Examiner: Kawsar; Abdullah Al

Assistant Examiner: Jablon; Asher H.

Attorney, Agent or Firm: Fish & Richardson P.C.

Background/Summary

BACKGROUND

(1) This specification relates to the characterization of activity in a recurrent artificial neural network. The characterization of activity can be used, e.g., in the identification of decision moments, as well as in encoding/decoding signals in contexts such as transmission, encryption, and data storage.

(2) Artificial neural networks are devices that are inspired by the structure and functional aspects of networks of biological neurons. In particular, artificial neural networks mimic the information encoding and other processing capabilities of networks of biological neurons using a system of interconnected constructs called nodes. The arrangement and strength of connections between nodes in an artificial neural network determines the results of information processing or information storage by the artificial neural network.

(3) Neural networks can be trained to produce a desired signal flow within the network and achieve desired information processing or information storage results. In general, training a neural network will change the arrangement and/or strength of connections between nodes during a learning phase. A neural network can be considered trained when sufficiently appropriate processing results are achieved by the neural network for given sets of inputs.

(4) Artificial neural networks can be used in a variety of different devices to perform non-linear data processing and analysis. Non-linear data processing does not satisfy the superposition principle, i.e., the variables that are to be determined cannot be written as a linear sum of independent components. Examples of contexts in which non-linear data processing is useful include pattern and sequence recognition, speech processing, novelty detection and sequential decision making, complex system modelling, and systems and techniques in a variety of other

contexts.

SUMMARY

(5) This specification describes technologies relating to the characterization of activity in an artificial neural network.

(6) For example, a method for identifying decision moments in a neural network includes determining a complexity of patterns of activity in the recurrent artificial neural network, wherein the activity is responsive to input into the recurrent artificial neural network, determining a timing of activity having a complexity that is distinguishable from other activity that is responsive to the input, and identifying the decision moments based on the timing of the activity that has the distinguishable complexity.

(7) As another example, a method for characterizing activity in a recurrent artificial neural network includes identifying clique patterns of activity of the recurrent artificial neural network. The method is performed by data processing apparatus.

(8) As another example, a method can include outputting a binary sequence of zeros and ones from a recurrent artificial neural network, wherein each digit in the sequence represents whether or not a particular group of nodes in the recurrent artificial neural network displays a respective pattern of activity.

(9) As another example, a method of structuring a recurrent artificial neural network can include characterizing a complexity of patterns of activity that can arise in the recurrent artificial neural network that comprises a structured collection of nodes and links between the nodes, and evolving the structure of the recurrent artificial neural network to increase the complexity of the patterns of activity. This method of structuring can also be used, e.g., a part of a method of training the recurrent artificial neural network.

(10) Other embodiments of these aspects include corresponding systems, apparatus, and computer programs, configured to perform the actions of the methods, encoded on computer storage devices.

(11) Particular embodiments of the subject matter described in this specification can be implemented so as to realize one or more of the following advantages. For example, traditional data processing devices such as, e.g., digital and other computers are programmed to follow a pre-defined logic sequence when processing information. The moment when a computer arrives at a result is thus relatively easy to identify. Namely, the completion of the logic sequence embodied in the programming indicates when information processing is complete and the computer has “reached a decision.” The result can be held at the output of the computer's data processor in a relatively long-lived form by, e.g., a memory device, a set of buffers, or the like, and accessed for a variety of purposes.

(12) In contrast, as described herein, decision moments in artificial recurrent neural networks can be identified based on characteristics of the dynamic properties of the neural network during information processing. Rather than waiting for the artificial neural network to reach a pre-defined end of a logic sequence, decision moments in artificial neural networks can be identified based on characteristics of the functional states of the artificial neural network in the midst of information processing.

(13) Further, the characteristics of the dynamic properties of recurrent artificial neural networks during information processing—including characteristics such as the activity comporting with clique patterns and directed clique patterns—can be used in a variety of signalling operations, including signal transmission, encoding, encryption, and storage. In particular, the characteristics of the activity in a recurrent artificial neural network during information processing reflects the input and can be considered to be an encoded form of the input (i.e., the “output” of a recurrent artificial neural network in encoding processes). These characteristics can be, e.g., transmitted to a remote receiver which could decode the transmitted characteristics to reconstitute the input or a portion of the input.

(14) Moreover, in some cases, activity in different groups of nodes a recurrent artificial neural

network (e.g., activity that comports with clique patterns and directed clique patterns) can be represented as a binary sequence of zeros and ones, each digit indicating whether or not that activity comports with a pattern. Since the activity can be the output of a recurrent artificial neural network in some contexts, the output of a recurrent artificial neural network can be represented as a vector of binary digits and compatible with digital data processing.

(15) Further, in some cases, such characterization of the dynamic properties of a recurrent artificial neural network can be used prior to and/or during training to increase the likelihood of complex patterns of activity arising during information processing. For example, before or during training, the links between nodes in a recurrent neural network can be intentionally evolved to increase the complexity of activity patterns. For example, the links between nodes in a recurrent artificial neural network can be intentionally evolved to increase the likelihood that, e.g., clique patterns and directed clique patterns of activity arise during information processing. This can reduce the time and effort needed to train a recurrent artificial neural network.

(16) As another example, such characterization of the dynamic properties of a recurrent artificial neural network can be used to determine a degree of completeness in the training of a recurrent neural network. For example, a recurrent artificial neural network that displays particular types of ordering in the activity (e. g, clique patterns and directed clique patterns) can be considered more highly trained than a recurrent artificial neural network that does not display such ordering. Indeed, in some cases, the extent of training can be quantified by quantifying the degree of ordering of the activity in the recurrent artificial neural network.

(17) The details of one or more implementations described in this specification are set forth in the accompanying drawings and the description below. Other features, aspects, and advantages of the subject matter will become apparent from the description, the drawings, and the claims.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a schematic illustration of the structure of recurrent artificial neural network device.

(2) FIGS. 2 and 3 are schematic illustrations of the function of recurrent artificial neural network device in different windows of time.

(3) FIG. 4 is a flowchart of a process for identifying decision moments in a recurrent artificial neural network based on characterization of the activity in the network.

(4) FIG. 5 is a schematic illustration of patterns of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network.

(5) FIG. 6 is a schematic illustration of patterns of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network.

(6) FIG. 7 is a schematic illustration of patterns of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network.

(7) FIG. 8 is a schematic illustration of a data table that can be used in a determination of the complexity or degree of ordering in the activity patterns in a recurrent artificial neural network device.

(8) FIG. 9 is a schematic illustration of a determination of the timing of activity patterns that have a distinguishable complexity.

(9) FIG. 10 is a flowchart of a process for encoding signals using a recurrent artificial neural network based on characterization of the activity in the network.

(10) FIG. 11 is a flowchart of a process for decoding signals using a recurrent artificial neural network based on characterization of the activity in the network.

(11) FIGS. 12, 13, and 14 are schematic illustrations of an identical binary form or representation of topological structures.

(12) FIGS. 15 and 16 schematically illustrate an example of how the presence or absence of features

that correspond to different bits are not independent of one another.

(13) FIGS. **17, 18, 19, 20** are schematic illustrations of the use of representations of the occurrence of topological structures in the activity in a neural network in four different classification systems.

(14) FIGS. **21, 22** are schematic illustration of edge devices that include a local artificial neural network that can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network.

(15) FIG. **23** is a schematic representation of a system in which local neural networks can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network.

(16) FIGS. **24, 25, 26, 27** are schematic illustrations of the use of representations of the occurrence of topological structures in the activity in a neural network in four different systems.

(17) FIG. **28** is a schematic illustration of a system **0** that includes an artificial neural network that can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network.

(18) Like reference symbols in the various drawings indicate like elements.

DETAILED DESCRIPTION

(19) FIG. **1** is a schematic illustration of the structure of a recurrent artificial neural network device **100**. Recurrent artificial neural network device **100** is a device that mimics the information encoding and other processing capabilities of networks of biological neurons using a system of interconnected nodes. Recurrent artificial neural network device **100** can be implemented in hardware, in software, or in combinations thereof.

(20) The illustration of recurrent artificial neural network device **100** includes a plurality of nodes **101, 102, . . . , 107** that are interconnected by a plurality of structural links **110**. Nodes **101, 102, . . . , 107** are discrete information processing constructs that are analogous to neurons in biological networks. Nodes **101, 102, . . . , 107** generally process one or more input signals received over one or more of links **110** to produce one or more output signals that are output over one or more of links **110**. For example, in some implementations, nodes **101, 102, . . . , 107** can be artificial neurons that weight and sum multiple input signals, pass the sum through one or more non-linear activation functions, and output one or more output signals.

(21) Nodes **101, 102, . . . , 107** can operate as accumulators. For example, nodes **101, 102, . . . , 107** can operate in accordance with an integrate-and-fire model in which one or more signals accumulate in a first node until a threshold is reached. After the threshold is reached, the first node fires by transmitting an output signal to a connected, second node along one or more of links **110**. In turn, the second node **101, 102, . . . , 107** accumulates the received signal and, if a threshold is reached, then the second node **101, 102, . . . , 107** transmits yet another output signal to a further connected node.

(22) Structural links **110** are connections that are capable of transmitting signals between nodes **101, 102, . . . , 107**. For the sake of convenience, all structural links **110** are treated herein as identical bidirectional links that convey a signal from every first of nodes **101, 102, . . . , 107** to every second of nodes **101, 102, . . . , 107** in identically the same manner as a signal is conveyed from the second to the first. However, this is not necessarily the case. For example, some portion or all of structural links **110** can be unidirectional links that convey a signal from a first of nodes **101, 102, . . . , 107** to a second of nodes **101, 102, . . . , 107** without conveying signals from the second to the first.

(23) As another example, in some implementations, structural links **110** can have diverse properties other than or in addition to directionality. For example, in some implementations, different structural links **110** can carry signals of different magnitudes—resulting in a different strengths of interconnection between respective of nodes **101, 102, . . . , 107**. As another example, different structural links **110** can carry different types of signal (e.g., inhibitory and/or excitatory signals). Indeed, in some implementations, structural links **110** can be modelled on the links between soma

in biological systems and reflect at least a portion of the enormous morphological, chemical, and other diversity of such links.

(24) In the illustrated implementation, recurrent artificial neural network device **100** is a clique network (or subnetwork) in that every node **101**, **102**, . . . , **107** is connected to every other node **101**, **102**, . . . , **107**. This is not necessarily the case. Rather, in some implementations, each node **101**, **102**, . . . , **107** can be connected to a proper subset of nodes **101**, **102**, . . . , **107** (by identical links or diverse links, as the case may be).

(25) For the sake of clarity of illustration, recurrent artificial neural network device **100** is illustrated with only seven nodes. In general, real-world neural network devices will include significantly larger numbers of nodes. For example, in some implementations, neural network devices can include hundreds of thousands, millions, or even billions of nodes. Thus, recurrent neural network device **100** can be a fraction of a larger recurrent artificial neural network (i.e., a subnetwork).

(26) In biological neural network devices, accumulation and signal transmission processes require the passage of time in the real world. For example, the soma of a neuron integrates input received over time, and signal transmission from neuron to neuron requires times that are determined by, e.g., the signal transmission velocity and the nature and length of the links between neurons. Thus, the state of a biological neural network device is dynamic and changes over time.

(27) In artificial recurrent neural network devices, time is artificial and represented using mathematical constructs. For example, rather than requiring a real world passage of time for signals to transmit from node to node, such signals can be represented in terms of artificial units that are generally unrelated to the real world passage of time—as measured in computer clock cycles or otherwise. Nevertheless, the state of an artificial recurrent neural network device can be described as “dynamic” in that it changes with respect to these artificial units.

(28) Please note that, for the sake of convenience, these artificial units are referred to herein as “time” units. Nevertheless, it is to be understood that these units are artificial and generally do not correspond to the real world passage of time.

(29) FIGS. 2 and 3 are schematic illustrations of the function of recurrent artificial neural network device **100** in different windows of time. Because the state of device **100** is dynamic, the functioning of device **100** can be represented using the signal transmission activity that occurs within a window. Such a functional illustration generally shows activity in only a fraction of links **110**. In particular, since in general not every link **110** conveys a signal within a particular window, not every link **110** is illustrated as actively contributing to the functioning of the device **100** in these illustrations.

(30) In the illustrations of FIGS. 2 and 3, an active link **110** is illustrated as a relatively thick solid line connecting a pair of nodes **101**, **102**, . . . , **107**. In contrast, inactive links **110** are illustrated as dashed lines. This is for the sake of illustration only. In other words, the structural connections formed by links **110** exist whether or not links **110** are active. However, this formalism highlights activity and the functioning of device **100**.

(31) In addition to schematically illustrating the existence of activity along a link, the direction of the activity is also schematically illustrated. In particular, the relatively thick solid lines that illustrate active links **110** also include arrow heads that denote the direction of signal transmission along the link during the relevant window. In general, the direction of signal transmission in a single window does not conclusively constrain the link to being a unidirectional link having the indicated directionality. Rather, in a first functional illustration for a first window of time, a link can be active in a first direction. In a second functional illustration for a second window, a link can be active in the opposite direction. However, in some cases such as, e.g., in a recurrent artificial neural network device **100** that exclusively includes unidirectional links, the directionality of signal transmission will conclusively indicate the directionality of the link.

(32) In feedforward neural network devices, information moves exclusively in a single direction

(i.e., forward) to an output layer of nodes that is at the end of the network. Feedforward neural network devices indicate that a “decision” has been reached and that information processing is complete by the propagation of the signals through the network to the output layer.

(33) In contrast, in recurrent neural networks, the connections between nodes form cycles and the activity of the network dynamically progresses without a readily identifiable decision. For example, even in a three-node recurrent neural network, the first node can transmit a signal to the second node, which in response can transmit a signal to the third. In response, the third node can transmit a signal back to the first. The signals received by the first node can be responsive—at least in part—to the signals transmitted from that same node.

(34) The schematic functional illustrations FIGS. 2 and 3 illustrate this in a network that is only slightly larger than a three-node recurrent neural network. The functional illustrations shown in FIG. 2 can be illustrative of activity within a first window and FIG. 3 can be illustrative of activity within a second, immediately following. As shown, a collection of signal transmission activity appears to originate in node **104** and progress in a generally clockwise direction through device **100** during the first window. In the second window, at least some of the signal transmission activity generally appears to return to node **104**. Even in such a simplistic illustration, signal transmission does not proceed in a manner that yields a clearly identifiable output or end.

(35) When a recurrent neural network of, e.g., thousands of nodes or more is considered, it can be recognized that signal propagation can occur over a huge number of paths and that these signals lack a clearly identifiable “output” location or time. Although the network may by design return to a quiescent state in which only background or even no signal transmission activity occurs, the quiescent state itself does not indicate the results of information processing. The recurrent neural network always returns to the quiescent state regardless of the input. The “output” or the result of the information processing is thus encoded in the activity that occurs within the recurrent neural network in response to a particular input.

(36) FIG. 4 is a flowchart of a process **400** for identifying decision moments in a recurrent artificial neural network based on characterization of the activity in the network. A decision moment is a point of time when the activity in a recurrent artificial neural network is indicative of the results of information processing by the network in response to an input. Process **400** can be performed by a system of one or more data processing apparatus that perform operations in accordance with the logic of one or more sets of machine-readable instructions. For example, process **400** can be performed by that same system of one or more computers that executes software for implementing the recurrent artificial neural network used in process **400**.

(37) The system performing process **400** receives a notification that a signal has been input into the recurrent artificial neural network at **405**. In some cases, the input of the signal is a discrete injection event in which, e.g., information is injected into one or more nodes and/or one or more links of the neural network. In other cases, the input of the signal is a stream of information that is injected into the one or more nodes and/or links of the neural network over a period of time. The notification indicates that the artificial neural network is actively processing information and not, e.g., in a quiescent state. In some cases, the notification is received from the neural network itself, e.g., such as when the neural network exits an identifiable quiescent state.

(38) The system performing process **400** divides the responsive activity in the network into a collection of windows at **410**. In cases where injection is a discrete event, the windows can subdivide the time between injection and a return to a quiescent state into a number of periods during which the activity displays variable complexities. In cases where injection is a stream of information, the duration of the injection (and optionally the time to return to a quiescent state after injection is complete) can be subdivided into windows during which the activity displays variable complexities. Various approaches to determining the complexity of activity are discussed further below.

(39) In some implementations, the windows all have the same duration, but this is not necessarily

the case. Rather, in some implementations, the windows can have different durations. For example, in some implementations, duration can increase as time since a discrete injection event has occurred increases.

(40) In some implementations, the windows can be a successive series of separate windows. In other implementations, the windows overlap in time so that one window begins before a previous window ends. In some cases, the windows can be a moving window that moves in time.

(41) In some implementations, different durations of windows are defined for different determinations of the complexity of activity. For example, for activity patterns that define activity occurring between relatively larger numbers of nodes, the windows can have a relatively longer duration than windows that are defined for activity patterns that define activity occurring between relatively smaller numbers of nodes. For example, in the context of patterns **500** of activity (FIG. 5), a window that is defined for identifying activity that comports with pattern **530** can be longer than a window that is defined for identifying activity that comports with pattern **505**.

(42) The system performing process **400** identifies patterns in the activity in the network in the different windows at **415**. As discussed further below, patterns in the activity can be identified by treating a functional graph as a topological space with nodes as points. In some implementation, the activity patterns that are identified are cliques, e.g., directed cliques, in a functional graph of the network.

(43) The system performing process **400** determines the complexity of the activity patterns in different windows at **420**. The complexity can be a measure of the likelihood that an ordered pattern of activity arises within a window. Thus, activity patterns that arise randomly would be relatively simple. On the other hand, activity patterns that show non-random order are relatively complex. For example, in some implementations, the complexity of an activity pattern can be measured using, e.g., the simplex counts or the Betti numbers of the activity pattern.

(44) The system performing process **400** determines the timing of activity patterns having a distinguishable complexity at **425**. A particular activity pattern can be distinguishable based on a complexity that deviates upwards or deviates downward, e.g., from a fixed or a variable baseline. In other words, the timing of activity patterns that indicate particularly high levels or particularly low levels of non-random order in the activity can be determined.

(45) For example, in cases where signal input is a discrete injection event, deviations, e.g., from a stable baseline or from a curve that is characteristic of the neural network's average response to a variety of different discrete injection events can be used to determine the timing of distinguishable complex activity patterns. As another example, in cases where information is input in a stream, large changes in complexity during streaming can be used to determine the timing of distinguishable complex activity patterns.

(46) The system performing process **400** times the reading of the output from the neural network based on the timing of distinguishably complex activity patterns at **430**. For example, in some implementations, the output of the neural network can be read at the same time that distinguishable complex activity patterns arise. In implementations where the complexity deviations indicate a relatively high non-random order in the activity, the observed activity patterns themselves can also be taken as the output of the recurrent artificial neural network.

(47) FIG. 5 is an illustration of patterns **500** of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network. For example, patterns **500** can be identified at **415** in process **400** (FIG. 4).

(48) Patterns **500** are illustrations of activity within a recurrent artificial neural network. During the application of patterns **500**, a functional graph is treated as a topological space with nodes as points. Activity in nodes and links that comports with patterns **500** can be recognized as ordered regardless of the identity of the particular nodes and/or links that participate in the activity. For example, first pattern **505** can represent the activity between nodes **101**, **104**, **105** in FIG. 2, with point 0 in pattern **505** as node **104**, point 1 as node **105**, and point 2 as node **101**. As another example, first pattern

505 can also represent the activity between nodes **104**, **105**, **106** in FIG. 3, with point 0 in pattern **505** as node **106**, point 1 as node **104**, and point 2 as node **105**. The order of activity in the directed cliques is also specified. For example, in pattern **505**, activity between point 1 and point 2 occurs after the activity between point 0 and point 1.

(49) In the illustrated implementation, patterns **500** are all directed cliques or directed simplices. In such patterns, activity originates from a source node that transmits signals to every other node in the pattern. In patterns **500**, such source nodes are designated as point 0 whereas the other nodes are designated as points 1, 2, Further, in directed cliques or simplices, one of the nodes acts a sink and receives signals transmitted from every other node in the pattern. In patterns **500**, such sink nodes are designated as the highest numbered point in the pattern. For example, in pattern **505**, the sink node is designated as point 2. In pattern **510**, the sink node is designated as point 3. In pattern **515**, the sink node is designated as point 3, and so on. The activity represented by patterns **500** is thus ordered in a distinguishable manner.

(50) Each of patterns **500** has a different number of points and reflects ordered activity in a different number of nodes. For example, pattern **505** is a 2D-simplex and reflects activity in three nodes, pattern **510** is a 3D-simplex and reflects activity in four nodes, and so on. As the number of points in a pattern increases, so does the degree of ordering and the complexity of the activity. For example, for a large collection of nodes that have a certain level of random activity within a window, some of that activity may comport with pattern **505** out of happenstance. However, it is progressively more unlikely that random activity will comport with the respective of patterns **510**, **515**, **520** The presence of activity that comports with pattern **530** indicates a relatively higher degree of ordering and complexity in the activity that the presence of activity that comports with pattern **505**.

(51) As discussed previously, in some implementations, different duration windows can be defined for different determinations of the complexity of activity. For example, when activity that comports with pattern **530** is to be identified, longer duration windows can be used than when activity that comports with pattern **505** is to be identified.

(52) FIG. 6 is an illustration of patterns **600** of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network. For example, patterns **600** can be identified at **415** in process **400** (FIG. 4).

(53) Like patterns **500**, patterns **600** are illustrations of activity within a recurrent artificial neural network. However, patterns **600** depart from the strict ordering of patterns **500** in that patterns **600** are not all directed cliques or directed simplices. In particular, patterns **605**, **610** have a lower directionality than pattern **515**. Indeed, pattern **605** lacks a sink node altogether. Nevertheless, patterns **605**, **610** indicate degree of ordered activity that exceeds that expected through random happenstance and can be used to determine the complexity of activity in a recurrent artificial neural network.

(54) FIG. 7 is an illustration of patterns **700** of activity that can be identified and used for identifying decision moments in a recurrent artificial neural network. For example, patterns **700** can be identified at **415** in process **400** (FIG. 4).

(55) Patterns **700** are groups of directed cliques or directed simplices of the same dimension (i.e., have the same number of points) that define patterns involving more points than the individual cliques or simplices and enclose cavities within the group of directed simplices.

(56) By way of example, pattern **705** includes six different three point, 2-dimensions patterns **505** that together define a homology class of degree two, whereas pattern **710** includes eight different three point, 2-dimensions patterns **505** that together define a second homology class of degree two. Each of the three point, 2-dimensions patterns **505** in patterns **705**, **710** can be thought of as enclosing a respective cavity. The n.sup.th Betti number associated with a directed graph provides a count of such homology classes within a topological representation.

(57) The activity illustrated by patterns such as patterns **700** illustrates a relatively high degree of

ordering of the activity within a network that is unlikely to arise by random happenstance. Patterns **700** can be used to characterize the complexity of that activity.

(58) In some implementations, only some patterns of activity are identified and/or some portion of the patterns of activity that are identified are discarded or otherwise ignored during the identification of decision moments. For example, with reference to FIG. 5, activity that comports with the five point, 4-dimensional simplex pattern **515** inherently includes activity that comports with the four point, 3-dimensional and three point, 2-dimension simplex patterns **510**, **505**. For example, points 0, 2, 3, 4 and points 1, 2, 3, 4 in 4-dimensional simplex pattern **515** of FIG. 5 both comport with 3-dimensional simplex pattern **510**. In some implementations, patterns that include fewer points—and hence are of a lower dimension—can be discarded or otherwise ignored during the identification of decision moments.

(59) As another example, only some patterns of activity need be identified. For example, in some implementations only patterns with odd number of points (3, 5, 7, . . .) or even numbers of dimensions (2, 4, 6, . . .) are used in the identification of decision moments.

(60) The complexity or degree of ordering in the activity patterns in a recurrent artificial neural network device in different windows can be determined in a variety of different ways. FIG. 8 is a schematic illustration of a data table **800** that can be used in such a determination. Data table **800** can be used to determine the complexity of the activity patterns in isolation or in conjunction with other activities. For example, data table **800** can be used at **420** in process **400** (FIG. 4).

(61) In further detail, table **800** includes a number count of pattern occurrences during a window “N,” where the number counts of activity that matches patterns of different dimensions are presented in different rows. For example, in the illustrated example, row **805** includes a number count (i.e., “2032”) of the occurrences of activity that matches one or more three point, 2-dimensional patterns, whereas row **810** includes a number count (i.e., “877”) of the occurrences of activity that matches one or more four point, 3-dimensional patterns. Since an occurrence of the patterns indicates that the activity has an order which is non-random, the number counts also provide generalized characterization of the overall complexity of the activity patterns. A table that is analogous to table **800** can be formed for each window that is defined, e.g., at **410** in process **400** (FIG. 4).

(62) Although table **800** includes a separate row **805**, **810**, **815**, **820**, **825** and a separate entry for every type of activity pattern, this is not necessarily the case. For example, one of more counts can be omitted (e.g., counts of simpler patterns) can be omitted from table **800** and from a determination of the complexity. As another example, some implementations, a single row or entry can include counts of occurrences of multiple activity patterns.

(63) Although FIG. 8 presents the number count in a table **800**, this is not necessarily the case. For example, the number count can be presented as a vector (e.g., <2032, 877, 133, 66, 48, . . . >). Regardless of how the count is presented, in some implementations, the counts can be expressed in binary and can be compatible with digital data processing infrastructure.

(64) In some implementations, number counts of the occurrences of the patterns can be weighted or combined to determine the degree or complexity of the ordering, e.g., at **420** in process **400** (FIG. 4). For example, the Euler characteristic can provide an approximation of the complexity of the activity and is given by:

$$S_{\text{sub}.0} - S_{\text{sub}.1} + S_{\text{sub}.2} - S_{\text{sub}.3} + \dots \quad \text{EQUATION 1}$$

where $S_{\text{sub}.n}$ is number of occurrences of a pattern of n points (i.e., a pattern of dimensionality $n-1$). The patterns can be, e.g., the directed clique patterns **500** (FIG. 5).

(65) As another example of how number counts of the occurrences of the patterns can be weighted to determine the degree or complexity of the ordering, in some implementations pattern occurrences can be weighted based on the weights of the links that are active. In further detail, as discussed previously, the strength of connection between nodes in an artificial neural network can vary, e.g., as a consequence of how active the connection was during training. An occurrence of a pattern of

activity along a collection of relatively stronger links can be weighted differently from the occurrence of that same pattern of activity along a collection of relatively weaker links. For example, in some implementations, the sum of the weight of the links that are active can be used to weight the occurrence.

(66) In some implementations, the Euler characteristic or other measure of complexity can be normalized by the total number of patterns that are matched within a particular window and/or the total number of patterns that it is possible for a given network to form given its structure. An example of a normalization with regard to the total number of patterns that it is possible for a network to form is given below in Equations 2, 3.

(67) In some implementations, occurrences of higher dimension patterns involving larger numbers of nodes can be weighted more heavily than occurrences of lower dimension patterns involving smaller numbers of nodes. For example, the probability of forming directed cliques decreases rapidly with increasing dimension. In particular, to form an n-clique from n+1 nodes, one needs (n+1)n/2 edges all oriented correctly. This probability can be reflected in the weighting.

(68) In some implementations, both the dimension and the directionality of patterns can be used to weight occurrences of patterns and determine the complexity of the activity. For example, with reference to FIG. 6, occurrences of five point, 4-dimensional pattern **515** can be weighted more heavily than occurrences of five point, 4-dimensional patterns **605**, **610** in accordance with the differences in directionality of those patterns.

(69) An example of the use of both the directionality and the dimension of patterns to determine the degree of ordering or complexity of the activity can be given by

$$(70) \quad \frac{S_0^{\text{active}}(n_0)^3 - S_1^{\text{active}}(n_1)^3 + S_2^{\text{active}}(n_2)^3 - S_3^{\text{active}}(n_3)^3 \cdot \text{Math.}}{\text{ERN}} / \text{SC} \quad \text{EQUATION 2}$$

where S.sub.x.sup.active indicates the number of active occurrences of a pattern of n points and ERN is the calculation for an equivalent random network, i.e., a network with the same number of nodes randomly connected. Further, SC is given by

$$(71) \quad \text{SC} = \frac{S_0^{\text{silent}}(n_0)^3 - S_1^{\text{silent}}(n_1)^3 + S_2^{\text{silent}}(n_2)^3 - S_3^{\text{silent}}(n_3)^3 \cdot \text{Math.}}{\text{ERN}} \quad \text{EQUATION 3}$$

where S.sub.x.sup.silent indicates the number of occurrences of a pattern of n points when the recurrent artificial neural network is silent and can be thought of as embodying the total number of patterns that it is possible for the network to form. In Equations 2, 3, the patterns can be, e.g., the directed clique patterns **500** (FIG. 5).

(72) FIG. 9 is a schematic illustration of a determination of the timing of activity patterns that have a distinguishable complexity. The determination illustrated in FIG. 9 can be performed in isolation or in conjunction with other activities. For example, the determination can be performed at **425** in process **400** (FIG. 4).

(73) FIG. 9 includes a graph **905** and a graph **910**. Graph **905** illustrates occurrences of patterns as a function of time along the x-axis. In particular, individual occurrences are illustrated schematically as vertical lines **906**, **907**, **908**, **909**. Each row of occurrences can be instances where activity matches a respective pattern or class of pattern. For example, the top row of occurrences can be instances where activity matches pattern **505** (FIG. 5), the second row of occurrences can be instances where activity matches pattern **510** (FIG. 5), the third row of occurrences can be instances where activity matches pattern **515** (FIG. 5), and so on.

(74) Graph **905** also includes dashed rectangles **915**, **920**, **925** that schematically delineate different windows of time when the activity patterns have a distinguishable complexity. As shown, the likelihood that activity in the recurrent artificial neural network matches a pattern indicative of complexity is higher during the windows delineated by dashed rectangles **915**, **920**, **925** than outside those windows.

(75) Graph **910** illustrates the complexity associated with these occurrences as a function of time

along the x-axis. Graph **910** includes a first peak **930** in complexity that coincides with the window delineated by dashed rectangle **915** and a second peak **935** in complexity that coincides with the window delineated by dashed rectangles **920**, **925**. As shown, the complexity illustrated by peaks **930**, **925** is distinguishable from what can be considered to be a baseline level **940** of complexity. (76) In some implementations, the times at which the output of a recurrent artificial neural network is to be read coincide with the occurrences of activity patterns that have a distinguishable complexity. For example, in the illustrative context of FIG. **9**, the output of a recurrent artificial neural network can be read at peaks **930**, **925**, i.e., during the windows delineated by dashed rectangles **915**, **920**, **925**.

(77) The identification of distinguishable levels of complexity in a recurrent artificial neural network is particularly beneficial when the input is a stream of data. Examples of data streams include, e.g., video or audio data. Although data streams have a beginning, it is generally desirable to process information in the data stream that does not have a pre-defined relationship with the beginning of the data stream. By way of example, a neural network could perform object recognition such as, e.g., recognizing bicyclists in the vicinity of an automobile. Such a neural networks should be able to recognizing bicyclists regardless of when those bicyclists appear in the video stream, i.e., without regard to time since the beginning of the video. Continuing with this example, when a data stream is input into an object recognition neural network, any patterns of activity in the neural network will generally display a low or quiescent level of complexity. These low or quiescent level of complexity are displayed regardless of the continuous (or nearly continuous) input of streaming data into the neural network device. However, when an object of interest appears in the video stream, the complexity of the activity will become distinguishable and indicate the time at which an object is recognized in the video stream. Thus, the timing of a distinguishable level of complexity of the activity can also act as a yes/no output as to whether the data in the data stream satisfies certain criteria.

(78) In some implementations, not only the timing but also the content of the output of the recurrent artificial neural network is given by the activity patterns that have a distinguishable complexity. In particular, the identity and activity of the nodes that participate in activity that comports with the activity patterns can be considered the output of the recurrent artificial neural network. The identified activity patterns can thus illustrate the result of processing by the neural network, as well as the timing when this decision is to be read.

(79) The content of the decision can be expressed in a variety of different forms. For example, in some implementations and as discussed in further detail below, the content of the decision can be expressed as a binary vector or matrix of ones and zeros. Each digit can indicate whether or not a pattern of activity is present or not, e.g., for a pre-defined group of nodes and/or a predefined duration. In such implementations, the content of the decision is expressed in binary and can be compatible with traditional digital data processing infrastructure.

(80) FIG. **10** is a flowchart of a process **1000** for encoding signals using a recurrent artificial neural network based on characterization of the activity in the network. Signals can be encoded in a variety of different contexts such as, e.g., transmission, encryption, and data storage. Process **1000** can be performed by a system of one or more data processing apparatus that perform operations in accordance with the logic of one or more sets of machine-readable instructions. For example, process **1000** can be performed by that same system of one or more computers that executes software for implementing the recurrent artificial neural network used in process **1000**. In some instances, process **1000** can be performed by the same data processing apparatus that performs process **400**. In some instances, process **1000** can be performed by an encoder that includes a computer system, e.g., the encoder in a signal transmission system or the encoder of a data storage system. In some examples the encoder is a video encoder.

(81) The system performing process **1000** inputs a signal into a recurrent artificial neural network at **1005**. In some cases, the input of the signal is a discrete injection event. In others, the input signal is

streamed into the recurrent artificial neural network.

(82) The system performing process **1000** identifies one or more decision moments in the recurrent artificial neural network at **1010**. For example, the system can identify one or more decision moments by performing process **400** (FIG. 4).

(83) The system performing process **1000** reads the output of the recurrent artificial neural network at **1015**. As discussed above, in some implementations, the content of the output of the recurrent artificial neural network is the activity in the neural network that matches the patterns used to identify the decision point(s).

(84) In some implementations, individual “reader nodes” can be added to a neural network to identify occurrences of a particular pattern of activity at a particular collection of nodes and hence to read the output of the recurrent artificial neural network at **1015**. The reader nodes can fire if and only if the activity at a particular collection of nodes satisfies timing (and possibly magnitude, as well) criteria. For example, in order to read an occurrence of pattern **505** (FIG. 5) at nodes **104**, **105**, **106** (FIGS. 2, 3), a reader node could be connected to nodes **104**, **105**, **106** (or the links **110** between them). The reader node would itself only become active if a pattern of activity involving nodes **104**, **105**, **106** (or their links) occurred.

(85) The use of such reader nodes would eliminate the need to define windows of time for the recurrent artificial neural network as a whole. In particular, individual reader nodes can be connected to different nodes and/or numbers of nodes (or the links between them). The individual reader nodes can be set to have tailored responses (e.g., different decay times in an integrate-and-fire model) to identify different activity patterns. The system performing process **1000** transmits or stores the output of the recurrent artificial neural network at **1020**. The particular action performed at **1020** can reflect the context in which process **1000** is being used. For example, in contexts where secure or compressed communications are desired, the system performing process **1000** can transmit the output of the recurrent neural network to a receiver that has access to the same or to a similar recurrent neural network. As another example, in contexts where secure or compressed data storage is desired, the system performing process **1000** can record the output of the recurrent neural network in one or more machine-readable data storage devices for later access.

(86) In some implementations, the complete output of the recurrent neural network is not transmitted or stored. For example, in implementations where the content of the output of the recurrent neural network is the activity in the neural network that matches the patterns indicative of complexity in the activity, only activity that matches relatively more complex or higher dimensional activity may be transmitted or stored. By way of example, in reference to patterns **500** (FIG. 5), in some implementations only activity that matches patterns **515**, **520**, **525**, and **530** is transmitted or stored whereas activity that matches patterns **505**, **510** is ignored or discarded. In this way, a lossy process allows the volume of data that is transmitted or stored to be reduced at the cost of the completeness of the information being encoded.

(87) FIG. 11 is a flowchart of a process **1100** for decoding signals using a recurrent artificial neural network based on characterization of the activity in the network. Signals can be decoded in a variety of different contexts such as, e.g., signal reception, decryption, and reading data from storage. Process **1100** can be performed by a system of one or more data processing apparatus that perform operations in accordance with the logic of one or more sets of machine-readable instructions. For example, process **1100** can be performed by that same system of one or more computers that executes software for implementing the recurrent artificial neural network used in process **1100**. In some instances, process **1100** can be performed by the same data processing apparatus that performs process **400** and/or process **1000**. In some instances, process **1100** can be performed by, e.g., the decoder in a signal reception system or the decoder of a data storage system.

(88) The system performing process **1100** receives at least a portion of the output of a recurrent artificial neural network at **1105**. The particular action performed at **1105** can reflect the context in which process **1100** is being used. For example, the system performing process **1000** can receive a

transmitted signal that includes the output of the recurrent artificial neural network or read a machine-readable data storage device that stores the output of the recurrent artificial neural network.

(89) The system performing process **1100** reconstructs the input of the recurrent artificial neural network from the received output at **1110**. Reconstruction can proceed in a variety of different ways. For example, in some implementations, a second artificial neural network (recurrent or not) can be trained to reconstruct the input into the recurrent neural network from the output received at **1105**.

(90) As another example, in some implementations, a decoder that has been trained using machine learning (including but not limited to deep learning) can be trained to reconstruct the input into the recurrent neural network from the output received at **1105**.

(91) As yet another example, in some implementations, input into the same recurrent artificial neural network or into a similar recurrent artificial neural network can be iteratively permuted until the output of that recurrent artificial neural network matches, to some degree, the output received at **1105**.

(92) In some implementations, process **1100** can include receiving user input specifying an extent to which the input is to be reconstructed and, in response, adjust the reconstruction at **1110** accordingly. For example, the user input could specify that a complete reconstruction is not needed. In response, the system performing process **1100** adjusts the reconstruction. For example, in implementations where the content of the output of the recurrent neural network is the activity in the neural network that matches the patterns indicative of complexity in the activity, only the output that characterizes activity that matches relatively more complex or higher dimensional activity would be used to reconstruct the input. By way of example, in reference to patterns **500** (FIG. 5), in some implementations only activity that matches patterns **515**, **520**, **525**, and **530** could be used to reconstruct the input, whereas activity that matches patterns **505**, **510** could be ignored or discarded. In this way, a lossy reconstruction can proceed in selected circumstances.

(93) In some implementations, processes **1000**, **1100** can be used for peer-to-peer encrypted communications. In particular, both the sender (i.e., the encoder) and the receiver (i.e., the decoder) can be provided with the same recurrent artificial neural network. There are several ways in which the shared recurrent artificial neural network can be tailored to ensure that a third party cannot reverse-engineer it and decrypt the signal, including: the structure of the recurrent artificial neural network the functional settings of the recurrent artificial neural network, including node states and edge weights, the size (or dimension) of the patterns, and the fraction of patterns in each dimension. These parameters can be thought of multiple layers that together ensure transmission security. Further, in some implementations, the decision moment time points can be used as keys to decrypt the signal.

(94) Although processes **1000**, **1100** are presented in terms of encoding and decoding a single recurrent artificial neural network, processes **1000**, **1100** can also be applied in systems and processes that rely upon multiple recurrent artificial neural networks. These recurrent artificial neural networks can operate either in parallel or in series.

(95) As an example of series operation, the output of a first recurrent artificial neural network can serve as the input of a second recurrent artificial neural network. The resultant output of the second recurrent artificial neural network is a twice encoded (or twice encrypted) version of the input into the first recurrent artificial neural network. Such a series arrangement of recurrent artificial neural networks can be useful in circumstances where different parties have different levels of access to information, e.g., in medical record systems where patient identity information may not be accessible to a party that will be using and have access to the remainder of the medical record.

(96) As an example of parallel operation, the same information can be input into multiple, different recurrent artificial neural networks. The different outputs of those neural networks can be used, e.g., to ensure that the input can be reconstructed with high fidelity.

(97) Although many implementations have been described, various modifications may be made. For example, although the application implies generally that activity within a recurrent artificial neural network should match a pattern indicative of ordering, this is not necessarily the case. Rather, in some implementations, activity within a recurrent artificial neural network can comport with a pattern without necessarily displaying activity that matches the pattern. For example, an increase in the likelihood that a recurrent neural network is to display activity that would match a pattern can be treated as non-random ordering of the activity.

(98) As yet another example, in some implementations, different groups of patterns can be tailored for use in characterizing the activity in different recurrent artificial neural networks. The patterns can be tailored, e.g., according to the effectiveness of the patterns in characterizing the activity of the different recurrent artificial neural networks. Effectiveness can be quantified, e.g., based on the size of a table or vector that represents the occurrence counts of different patterns.

(99) As yet another example, in some implementations, the patterns used to characterize the activity in a recurrent artificial neural networks can consider the strength of a connection between nodes. In other words, the patterns described previously herein treat all signal transmission activity between two nodes in a binary manner, i.e., either the activity exists or it doesn't. This is not necessarily the case. Rather, in some implementations, comporting with a pattern can require activity of a certain level or strength of connection to be taken as indicative of ordered complexity in the activity of a recurrent artificial neural network.

(100) As yet another example, the content of the output of the recurrent artificial neural network can include activity patterns that occur outside windows of time in which the activity in a neural network has a distinguishable level of complexity. For example, the output of the recurrent artificial neural network that is read at **1015** and transmitted or stored at **1020** (FIG. **10**) can include information encoding activity patterns that occur, e.g., outside dashed rectangles **915**, **920**, **925** in graph **905** (FIG. **9**). By way of example, the output of the recurrent artificial neural network could characterize only the highest dimensional patterns of activity, regardless of when those patterns of activity occur. As another example, the output of the recurrent artificial neural network could characterize only patterns of activity that enclose cavities, regardless of when those patterns of activity occur.

(101) FIGS. **12**, **13**, and **14** are schematic illustrations of an identical binary form or representation **1200** of topological structures such as, e.g., patterns of activity in a neural network. The topological structures illustrated in FIGS. **12**, **13**, and **14** all include the same information, namely, an indication the presence or absence of features in a graph. The features can be, e.g., activity in a neural network device. In some implementations, the activity is identified based on or during periods of time in which the activity in the neural network has a complexity that is distinguishable from other activity that is responsive to an input.

(102) As shown, binary representation **1200** includes bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** and an additional, arbitrary number of bits (represented by the ellipses "..."). For didactic purposes, bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** ... are illustrated as discrete rectangular shapes that are either filled or unfilled to indicate the binary value of the bit. In the schematic illustrations, representation **1200** superficially appears to be either a one-dimensional vector of bits (FIGS. **12**, **13**) or a two-dimensional matrix of bits (FIG. **14**). However, representation **1200** differs from a vector, from a matrix, or from other ordered collection of bits in that the same information can be encoded regardless of the order of the bits—i.e., regardless of the location of individual bits within the collection.

(103) For example, in some implementations, each individual bit **1205**, **1207**, **1211**, **1293**, **1294**, **1297** ... can represent the presence or absence of a topological feature—regardless of the location of that feature in the graph. By way of example, referring to FIG. **12**, a bit such as bit **1207** can indicate the presence of a topological feature that comports with pattern **505** (FIG. **5**), regardless of whether that activity occurs between nodes **104**, **105**, **101** or between nodes **105**, **101**, **102**. Thus,

although each individual bit **1205**, **1207**, **1211**, **1293**, **1294**, **1297**, . . . can be associated with a particular feature, the location of that feature in the graph need not be encoded, e.g., by a corresponding location of the bit in representation **1200**. In other words, in some implementations, representation **1200** may only provide an isomorphic topological reconstruction of the graph. (104) As an aside, in other implementations, it is possible that the location of individual bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297**, . . . does indeed encode information such as, e.g., the location of a feature in the graph. In these implementations, the source graph can be reconstructed using representation **1200**. However, such an encoding is not necessarily present.

(105) In view of the ability of a bit to represent the presence or absence of a topological feature regardless of the location of that feature in the graph, in FIG. 12, bit **1205** appears before bit **1207**, which appears before bit **1211** at the start of representation **1200**. In contrast, in FIGS. 13 and 14, the order of bits **1205**, **1207**, and **1211** within representation **1200**—and the position of bits **1205**, **1207**, and **1211** relative to other bits within representation **1200**—has changed. Nevertheless, binary representation **1200** remains the same—as does the set of rules or algorithm that defines the process for encoding information in binary representation **1200**. So long as the correspondence between the bit and the feature is known, the location of the bits in the representation **1200** is irrelevant.

(106) In further detail, each bit **1205**, **1207**, **1211**, **1293**, **1294**, **1297** . . . individually represents the presence or absence of a feature in a graph. A graph is a set of nodes and a set of edges between those nodes. The nodes can correspond to objects. Examples of objects can include, e.g., artificial neurons in a neural network, individuals in a social network, or the like. Edges can correspond to some relation between the objects. Examples of relations include, e.g., a structural connection or activity along the connection. In the context of a neural network, artificial neurons can be related by a structural connection between neurons or by transmission of information along a structural connection. In the context of a social network, individuals can be related by a “friend” or other relational connection or by transmission of information (e.g., a posting) along such a connection. Edges can thus characterize relatively long-lived structural characteristics of the set of nodes or relatively transient activity characteristics that occur within a defined time frame. Further, edges can either be directed or bidirectional. Directed edges indicate directionality of the relation between the objects. For example, the transmission of information from a first neuron to a second neuron can be represented by a directed edge that denotes the direction of transmission. As another example, in a social network, a relational connection may indicate that second user is to receive information from the first but not that the first is to receive information from the second. In topological terms, a graph can be expressed as a set of unit intervals $[0, 1]$ where 0 and 1 are identified with respective nodes that are connected by an edge.

(107) The features whose presence or absence is indicated by bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** can be, e.g., a node, a set of nodes, a set of sets of nodes, a set of edges, a set of sets of edges, and/or additional hierarchically-more-complex features (e.g., a set of sets of sets of nodes). Bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** generally represent the presence or absence of features that are at different hierarchical levels. For example, bit **1205** may represent the presence or absence of a node, whereas bit **1207** may represent the presence or absence of a set of nodes.

(108) In some implementations, bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** may represent features in a graph that have a threshold level of some characteristic. For example, bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297** can represent not only that there is activity in a set of edges, but also that this activity is weighted either above or below a threshold level. The weights can, e.g., embody the training of a neural network device to a particular purpose or can be an innate characteristic of the edges.

(109) FIGS. 5, 6, and 8 above illustrate features whose presence or absence can be represented by bits **1205**, **1207**, **1211**, **1293**, **1294**, **1297**

(110) The directed simplices in collections **500**, **600**, **700** treat functional or structural graphs as a topological space with nodes as points. Structure or activity involving one or more nodes and links

that comports with simplices in collection **500, 600, 700** can be represented in a bit regardless of the identity of the particular nodes and/or links that participate in the activity.

(111) In some implementations, only some patterns of structure or activity are identified and/or some portion of the patterns of structure or activity that are identified are discarded or otherwise ignored. For example, with reference to FIG. 5, structure or activity that comports with the five point, 4-dimensional simplex pattern **515** inherently includes structure or activity that comports with the four point, 3-dimensional and three point, 2-dimension simplex patterns **510, 505**. For example, points 0, 2, 3, 4 and points 1, 2, 3, 4 in 4-dimensional simplex pattern **515** of FIG. 5 both comport with 3-dimensional simplex pattern **510**. In some implementations, simplex patterns that include fewer points—and hence are of a lower dimension—can be discarded or otherwise ignored.

(112) As another example, only some patterns of structure or activity need be identified. For example, in some implementations only patterns with odd number of points (3, 5, 7, . . .) or even numbers of dimensions (2, 4, 6, . . .) are used.

(113) Returning to FIGS. **12, 13, 14**, the features whose presence or absence is represented by bits **1205, 1207, 1211, 1293, 1294, 1297 . . .** may not be independent of one another. By way of explanation, if bits **1205, 1207, 1211, 1293, 1294, 1297** represent the presence or absence of 0D-simplices that each reflect the existence or activity of a single node, then bits **1205, 1207, 1211, 1293, 1294, 1297** are independent of one another. However, if bits **1205, 1207, 1211, 1293, 1294, 1297** represent the presence or absence of higher-dimensional simplices that each reflect the existence or activity of multiple nodes, then the information encoded by the presence or absence of each individual feature may not be independent of the presence or absence of the other features.

(114) FIG. **15** schematically illustrates an example of how the presence or absence of features that correspond to different bits are not independent of one another. In particular, a subgraph **1500** that includes four nodes **1505, 1510, 1515, 1520** and six directed edges **1525, 1530, 1535, 1540, 1545, 1550** is illustrated. In particular, edge **1525** is directed from node **1525** to node **1510**, edge **1530** is directed from node **1515** to node **1505**, edge **1535** is directed from node **1520** to node **1505**, edge **1540** is directed from node **1520** to node **1510**, edge **1545** is directed from node **1515** to node **1510**, edge **1550** is directed from node **1515** to node **1520**.

(115) A single bit in representation **1200** (e.g., filled bit **1207** in FIGS. **12, 13, 14**) may indicate the presence of a directed 3D-simplex. For example, such a bit could indicate the presence of the 3D-simplex formed by nodes **1505, 1510, 1515, 1520** and edges **1525, 1530, 1535, 1540, 1545, 1550**. A second bit in representation **1200** (e.g., filled bit **1293** in FIGS. **12, 13, 14**) may indicate the presence of a directed 2D-simplex. For example, such a bit could indicate the presence of the 2D-simplex formed by nodes **1515, 1505, 1510** and edges **1525, 1530, 1545**. In this simple example, the information encoded by bit **1293** is completely redundant with the information encoded by bit **1207**.

(116) Please note that the information encoded by bit **1293** may also be redundant with the information encoded by still further bit. For example, the information encoded by bit **1293** would be redundant with both a third bit and a fourth bit that indicates the presence of additional directed 2D-simplices. Examples of those simplices are formed by nodes **1515, 1520, 1510** and edges **1540, 1545, 1550** and nodes **1520, 1505, 1510** and edges **1525, 1535, 1540**.

(117) FIG. **16** schematically illustrates another example of how the presence or absence of features that correspond to different bits are not independent of one another. In particular, a subgraph **1600** that includes four nodes **1505, 1510, 1515, 1520** and five directed edges **1625, 1630, 1635, 1640, 1645** is illustrated. Nodes **1505, 1510, 1515, 1520** and edges **1625, 1630, 1635, 1640, 1645** generally correspond to nodes **1505, 1510, 1515, 1520** and edges **1525, 1530, 1535, 1540, 1545** in subgraph **1500** (FIG. **15**). However, in contrast with subgraph **1500** in which nodes **1515, 1520** are connected by edge **1550**, nodes **1615, 1620** are not connected by an edge.

(118) A single bit in representation **1200** (e.g., unfilled bit **1205** in FIGS. **12, 13, 14**) may indicate the absence of a directed 3D-simplex, such as, e.g., the directed 3D-simplex that encompasses

nodes **1605, 1610, 1615, 1620**. A second bit in representation **1200** (e.g., filled bit **1293** in FIGS. **12, 13, 14**) may indicate the presence of a 2D-simplex. An example directed 2D-simplex is formed by nodes **1615, 1605, 1610** and edges **1625, 1630, 1645**. This combination of a filled bit **1293** and an unfilled bit **1205** provides information indicative of presence or absence of other features (and the state of other bits) that may or may not be present in representation **1200**. In particular, the combination of the absence of a directed 3D-simplex and the presence of a directed 2D-simplex indicates that at least one edge is absent from either: a) the possible directed 2D-simplex formed by nodes **1615, 1620, 1610** or b) the possible directed 2D-simplex formed by nodes **1620, 1605, 1610**. (119) Thus, the state of a bit that represents the presence or absence of either of these possible simplices is not independent of the state of bits **1205, 1293**.

(120) Although these examples have been discussed in terms of a features with different numbers of nodes and a hierarchical relationship, this is not necessarily the case. For example, a representation **1200** that includes a collection of bits that corresponds only to, e.g., the presence or absence of 3D-simplices is possible.

(121) Using individual bits to represent the presence or absence of features in a graph yields certain properties. For example, the encoding of information is fault tolerant and provides “graceful degradation” of the encoded information. In particular, the loss of a particular bit (or group of bits) may increase the uncertainty as to the presence or absence of a feature. However, estimates of the likelihood that a feature is present or absent will still be possible from the other bits that indicate the presence or absence of adjacent features.

(122) Likewise, as the number of bits increases, certainty as to the presence or absence of a feature increases.

(123) As another example, as discussed above, the ordering or arrangement of bits is irrelevant to isomorphic reconstruction of the graph that is represented by the bit. All that is required is a known correspondence between the bits and particular nodes/structures in the graph.

(124) In some implementations, the patterns of activity in a neural network can be encoded in a representation **1200** (FIGS. **12, 13, and 14**). In general, the patterns of activity in a neural network are a result of a number of characteristics of the neural network such as, e.g., the structural connections between nodes of the neural network, the weights between nodes, as well as a whole host of possible other parameters. For example, in some implementations, the neural network could have been trained prior to the encoding of the patterns of activity in representation **1200**.

(125) However, regardless of whether the neural network untrained or trained, for a given input, the responsive pattern of activity can be thought of a “representation” or an “abstraction” of that input within the neural network. Thus, although representation **1200** can appear to be a straightforward-appearing collection of (in some cases, binary) digits, each of the digits can encode the relationship or correspondence between a particular input and relevant activity in the neural network.

(126) FIGS. **17, 18, 19, 20** are schematic illustrations of the use of representations of the occurrence of topological structures in the activity in a neural network in four different classification systems **1700, 1800, 1900, 2000**. Classification systems **1700, 1800** each classify representations of the patterns of activity in a neural network as part of the classification of input. Classification systems **1900, 2000** each classify approximations of representations of the patterns of activity in a neural network as part of the classification of input. In classification systems **1700, 1800**, the patterns of activity that are represented occur in and are read from a source neural network device **1705** that is part of the classification system **1700, 1800**. In contrast, in classification systems **1900, 2000**, the patterns of activity that are approximately represented occur in a source neural network device that is not part of the classification system **1700, 1800**. Nevertheless, the approximation of the representation of those patterns of activity are read from an approximator **1905** that is part of classification systems **1900, 2000**.

(127) In additional detail, turning to FIG. **17**, classification system **1700** includes a source neural network **1705** and a linear classifier **1710**. Source neural network **1705** is a neural network device

that is configured to receive an input and present representations of the occurrence of topological structures in the activity within source neural network **1705**. In the illustrated implementation, source neural network **1705** includes an input layer **1715** that receives the input. However, this is not necessarily the case. For example, in some implementation, some or all of the input can be injected into different layers and/or edges or nodes throughout source neural network **1705**.

(128) Source neural network **1705** can be any of variety of different types of neural network. In general, source neural network **1705** is a recurrent neural network such as, e.g., a recurrent neural network that is modelled on a biological system. In some cases, source neural network **1705** can model a degree of the morphological, chemical, and other characteristics of a biological system. In general, source neural network **1705** is implemented on one or more computing devices with a relatively high level of computational performance, e.g., a supercomputer. In such cases, classification system **1700** will generally be a dispersed system in which a remote classifier **1710** communicates with source neural network **1705**, e.g., via a data communications network.

(129) In some implementations, source neural network **1705** can be untrained and the activity that is represented can be the innate activity of source neural network **1705**. In other implementations, source neural network **1705** can be trained and the activity that is represented can embody this training.

(130) The representations read from source neural network **1705** can be representations such as representation **1200** (FIGS. **12**, **13**, **14**). The representations can be read from source neural network **1705** in a number of ways. For example, in the illustrated example, source neural network **1705** includes “reader nodes” that read patterns of activity between other nodes within source neural network **1705**. In other implementations, the activity within source neural network **1705** are read by a data processing component of that is programmed to monitor source neural network **1705** for relatively highly-ordered patterns of activity. In still other implementations, source neural network **1705** can include an output layer from which representation **1200** can be read, e.g., when source neural network **1705** is implemented as a feed-forward neural network.

(131) Linear classifier **1710** is a device that classifies an object—namely, representations of the patterns of activity in source neural network **1705**—based on a linear combination of the object's characteristics. Linear classifier **1710** includes an input **1720** and an output **1725**. Input **1720** is coupled to receive representations of the patterns of activity in source neural network **1705**. In other words, the representations of the patterns of activity in source neural network **1705** is a feature vector that represents the characteristics of the input into source neural network **1705** that are used by linear classifier **1710** to classify that input. Linear classifier **1710** can receive the representations of the patterns of activity in source neural network **1705** in a variety of ways. For example, the representations of the patterns of activity can be received as discrete events or as a continuous stream over a real time or non-real time communication channel.

(132) Output **1725** is coupled to output the classification result from linear classifier **1710**. In the illustrated implementation, output **1725** is schematically illustrated as a parallel port with multiple channels. This is not necessarily the case. For example, output **1725** can output classification result over a serial port or a port with combined parallel and serial capabilities.

(133) In some implementations, linear classifier **1710** can be implemented on one or more computing devices with relatively limited computational performance. For example, linear classifier **1710** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet.

(134) In FIG. **18**, classification system **1800** includes source neural network **1705** and a neural network classifier **1810**. Neural network classifier **1810** is a neural network device classifies an object—namely, representations of the patterns of activity in source neural network **1705**—based on a non-linear combination of the object's characteristics. In the illustrated implementation, neural network classifier **1810** is a feedforward network that includes an input layer **1820** and an output layer **1825**. As with linear classifier **1710**, neural network classifier **1810** can receive the

representations of the patterns of activity in source neural network **1705** in a variety of ways. For example, the representations of the patterns of activity can be received as discrete events or as a continuous stream over a real time or non-real time communication channel.

(135) In some implementations, neural network classifier **1810** can perform inferences on one or more computing devices with relatively limited computational performance. For example, neural network classifier **1810** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet, e.g., in a Neural Processing Unit of such a device. Like classification system **1700**, classification system **1800** will generally be a dispersed system in which a remote neural network classifier **1810** communicates with source neural network **1705**, e.g., via a data communications network.

(136) In some implementations, neural network classifier **1810** can be, e.g., a deep neural network such as a convolutional neural network that includes convolutional layers, pooling layers, and fully-connected layers. Convolutional layers can generate feature maps, e.g., using linear convolutional filters and/or nonlinear activation functions. Pooling layers reduce the number of parameters and control overfitting. The computations performed by the different layers in image classifier **1820** can be defined in different ways in different implementations of image classifier **1820**.

(137) In FIG. **19**, classification system **1900** includes source approximator **1905** and a linear classifier **1710**. As discussed further below, source approximator **1905** is a relatively simple neural network that is trained to receive input—either at an input layer **1915** or elsewhere—and output a vector that approximates a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network. For example, source approximator **1905** can be trained to approximate a recurrent source neural network such as, e.g., a recurrent neural network that is modelled on a biological system and includes a degree of the morphological, chemical, and other characteristics of a biological system. In the illustrated implementation, source approximator **1905** includes an input layer **1915** and an output layer **1920**. Input layer **1915** is coupleable to receive the input data. Output layer **1920** is coupled to output an approximation of a representation of the activity within a neural network device for receipt by input **1720** of linear classifier. For example, output layer **1920** can output an approximation **1200'** of representation **1200** (FIGS. **12**, **13**, **14**). As an aside, the representation **1200** schematically illustrated in FIGS. **17**, **18** and the approximation **1200'** of representation **1200** schematically illustrated in FIGS. **19**, **20** are identical. This is for the sake of convenience only. In general, approximation **1200'** will differ from representation **1200** in at least some ways. Despite those differences, linear classifier **1710** can classify approximation **1200'**.

(138) In general, source approximator **1905** can perform inferences on one or more computing devices with relatively limited computational performance. For example, source approximator **1905** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet, e.g., in a Neural Processing Unit of such a device. In general and in contrast with classification systems **1700**, **1800**, classification system **1900** will generally be housed in a single housing, e.g., with source approximator **1905** and linear classifier **1710** implemented on either the same data processing devices or on data processing devices coupled by a hardwired connection.

(139) In FIG. **20**, classification system **2000** includes source approximator **1905** and a neural network classifier **1810**. Output layer **1920** of source approximator **1905** is coupled to output an approximation **1200'** of a representation of the activity within a neural network device for receipt by input **1820** of neural network classifier **1810**. Despite any differences between approximation **1200'** and representation **1200**, neural network classifier **1810** can classify approximation **1200'**. In general and like classification system **1900**, classification system **2000** will generally be housed in a single housing, e.g., with source approximator **1905** and neural network classifier **1810** implemented on either the same data processing devices or on data processing devices coupled by a hardwired connection.

(140) FIG. **21** is a schematic illustration of an edge device **2100** that includes a local artificial neural network that can be trained using representations of the occurrence of topological structures

corresponding to activity in a source neural network. In this context, a local artificial neural network can be, e.g., an artificial neural network that is executed entirely on one or more local processors that do not require a communications network to exchange data. In general, the local processors will be connected by hardwire connections. In some instances, the local processors can be housed in a single housing, such as a single personal computer or a single handheld, mobile device. In some instances, the local processors can be under the control of and accessible by a single individual or a limited number of individuals. In effect, by using a representation of the occurrence of topological structures in a more complex source neural network to train (e.g., using a supervised learning or reinforcement learning technique) a simpler and/or less highly trained but more exclusive second neural network, even individuals with limited computing resources and limited numbers of training samples can train a neural network as desired. Storage requirements and computational complexity during training are reduced and resources like battery lifespan are spared.

(141) In the illustrated implementation, edge device **2100** is schematically illustrated as a security-camera device that includes an optical imaging system **2110**, image processing electronics **2115**, a source approximator **2120**, a representation classifier **2125**, and a communications controller and interface **2130**.

(142) Optical imaging system **2110** can include, e.g., one or more lenses (or even a pinhole) and a CCD device. Image processing electronics **2115** can read the output of optical imaging system **2110** and in general can perform basic image processing functions. Communications controller and interface **2130** is a device that is configured to control the flow of information to and from device **2100**. As discussed further below, among the operations that communications controller and interface **2130** can perform are the transmission of images of interest to other devices and the receipt of training information from other devices. Communications controller and interface **2130** can thus include both a data transmitter and a receiver that can communicate over, e.g., a data port **2135**. Data port **2135** can be a wired port, a wireless port, an optical port, or the like.

(143) Source approximator **2120** is relatively simple neural network that is trained to output a vector that approximates a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network. For example, source approximator **2120** can be trained to approximate a recurrent source neural network such as, e.g., a recurrent neural network that is modelled on a biological system and includes a degree of the morphological, chemical, and other characteristics of a biological system.

(144) Representation classifier **2125** is either a linear classifier or a neural network classifier that is coupled to receive an approximation of a representation of the patterns of activity in a source neural network from source approximator **2120** and output a classification result. Representation classifier **2125** can be, e.g., a deep neural network such as a convolutional neural network that includes convolutional layers, pooling layers, and fully-connected layers. Convolutional layers can generate feature maps, e.g., using linear convolutional filters and/or nonlinear activation functions. Pooling layers reduce the number of parameters and control overfitting. The computations performed by the different layers in representation classifier **2125** can be defined in different ways in different implementations of representation classifier **2125**.

(145) In some implementations, in operation, optical imaging system **2110** can generate raw digital images. Image processing electronics **2115** can read the raw images and will generally perform at least some basic image processing functions. Source approximator **2120** can receive images from image processing electronics **2115** and perform inference operations to output a vector that approximates a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network. This approximation vector is input into representation classifier **2125** which determines whether the approximation vector satisfies one or more sets of classification criteria. Examples include facial recognition and other machine vision operations. In the event that representation classifier **2125** determines that the approximation vector satisfies a set of classification criteria, representation classifier **2125** can instruct communications controller and

interface **2130** to transmit information regarding the images. For example, communications controller and interface **2130** can transmit the image itself, the classification, and/or other information regarding the images.

(146) At times, it may be desirable to change the classification process. In these cases, communications controller and interface **2130** can receive a training set. In some implementations, the training set can include raw or processed image data and representations of topological structures that arise in the patterns of activity in a relatively more complex neural network. Such a training set can be used to retrain source approximator **2120**, e.g., using a supervised learning or reinforcement learning technique. In particular, the representations are used as the target answer vectors and represent the desired result of source approximator **2120** processing the raw or processed image data.

(147) In other implementations, the training set can include representations of topological structures that arise in the patterns of activity in a relatively more complex neural network and the desired classification of those representations of topological structures. Such a training set can be used to retrain a neural network representation classifier **2125**, e.g., using a supervised learning or reinforcement learning technique. In particular, the desired classification are used as the target answer vectors and represent the desired result of representation classifier **2125** processing the representations of topological structures.

(148) Regardless of whether source approximator **2120** or representation classifier **2125** is retrained, inference operations at device **2100** can be facily adapted to changing circumstances and objectives without large sets of training data and time- and computing power-intensive iterative training.

(149) FIG. **22** is a schematic illustration of a second edge device **2200** that includes a local artificial neural network that can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network. In the illustrated implementation, second edge device **2200** is schematically illustrated as a mobile computing device such as smart phone or a tablet. Device **2200** includes an optical imaging system (e.g., on the backside of device **2200**, not shown), image processing electronics **2215**, a representation classifier **2225**, a communications controller and interface **2230**, and a data port **2235**. These components can have characteristics and perform actions that correspond to those of optical imaging system **2110**, image processing electronics **2115**, representation classifier **2125**, communications controller and interface **2130**, and data port **2135** in device **2100** (FIG. **21**).

(150) The illustrated implementation of device **2200** additionally includes one or more additional sensors **2240** and a multi-input source approximator **2245**. Sensor(s) **2240** can sense one of more characteristics of the environment surrounding device **2200** or of device **2200** itself. For example, in some implementations, sensor **2240** can be an accelerometer that senses the acceleration to which device **2200** is subject. As another example, in some implementations, sensor **2240** can be an acoustic sensor such as a microphone that senses noise in the environment of device **2200**. Still further examples of sensor **2240** include chemical sensors (e.g., “artificial noses” and the like), humidity sensors, radiation sensors, and the like. In some cases, sensor **2240** is coupled to processing electronics that can read the output of sensor **2240** (or other information such as, e.g., a contact list or map) and perform basic processing functions. Different implementations of sensor **2240** can thus have different “modalities” in that the physical sensed physical parameter changes from sensor to sensor.

(151) Multi-input source approximator **2245** is a relatively simple neural network that is trained to output a vector that approximates a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network. For example, multi-input source approximator **2245** can be trained to approximate a recurrent source neural network such as, e.g., a recurrent neural network that is modelled on a biological system and includes a degree of the morphological, chemical, and other characteristics of a biological system.

(152) Unlike source approximator **2120**, multi-input source approximator **2245** is coupled to receive raw or processed sensor data from multiple sensors and return an approximation of a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network based on that data. For example, multi-input source approximator **2245** can receive processed image data from image processing electronics **2215** as well as, e.g., acoustic, acceleration, chemical, or other data from one or more sensors **2240**. Multi-input source approximator **2245** can be, e.g., a deep neural network such as a convolutional neural network that includes convolutional layers, pooling layers, and fully-connected layers. The computations performed by the different layers in multi-input source approximator **2245** can be dedicated to a single type of sensor data or to sensor data of multiple modalities.

(153) Regardless of the particular organization of multi-input source approximator **2245**, multi-input source approximator **2245** is to return the approximation based on raw or processed sensor data from multiple sensors. For example, processed image data from image processing electronics **2215** and acoustic data from a microphone sensor **2240** can be used by multi-input source approximator **2245** to approximate a representation of topological structures that would arise in the patterns of activity in a relatively more complex neural network that received the same data.

(154) At times, it may be desirable to change the classification process at device **2200**. In these cases, communications controller and interface **2230** can receive a training set. In some implementations, the training set can include raw or processed image, sounds, chemical or other data and representations of topological structures that arise in the patterns of activity in a relatively more complex neural network. Such a training set can be used to retrain multi-input source approximator **2245**, e.g., using a supervised learning or reinforcement learning technique. In particular, the representations are used as the target answer vectors and represent the desired result of multi-input source approximator **2245** processing the raw or processed image or sensor data.

(155) In other implementations, the training set can include representations of topological structures that arise in the patterns of activity in a relatively more complex neural network and the desired classification of those representations of topological structures. Such a training set can be used to retrain a neural network representation classifier **2225**, e.g., using a supervised learning or reinforcement learning technique. In particular, the desired classification are used as the target answer vectors and represent the desired result of representation classifier **2225** processing the representations of topological structures.

(156) Regardless of whether multi-input source approximator **2245** or representation classifier **2225** is retrained, inference operations at device **2200** can be facily adapted to changing circumstances and objectives without large sets of training data and time- and computing power-intensive iterative training.

(157) FIG. **23** is a schematic representation of a system **2300** in which local neural networks can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network. The target neural networks are implemented on relatively simple, less expensive data processing systems whereas the source neural network can implemented on a relatively complex, more expensive data processing system.

(158) System **2300** includes a variety of devices **2305** with local neural networks, a telephone base station **2310**, a wireless access point **2315**, a server system **2320**, and one or more data communications networks **2325**.

(159) Local neural networks devices **2305** are devices that are configured to process data using computationally-less-intensive target neural networks. As illustrated, local neural networks devices **2305** can be implemented as mobile computing devices, cameras, automobiles, or any of a legion of other appliances, fixtures, and mobile components, as well as different makes and models of devices within each category. Different local neural networks devices **2305** can belong to different owners. In some implementations, access to the data processing functionality of local neural networks devices **2305** will generally be restricted to these owners and/or the owner's designates.

(160) Local neural networks devices **2305** can each include one or more source approximators that are trained to output a vector that approximates a representation of topological structures that arise in the patterns of activity in a relatively more complex neural network. For example, the relatively more complex neural network can be a recurrent source neural network such as, e.g., a recurrent neural network that is modelled on a biological system and includes a degree of the morphological, chemical, and other characteristics of a biological system.

(161) In some implementations, in addition to processing data using source approximators, local neural networks devices **2305** can also be programmed to re-train the source approximators using representations of topological structures that arise in the patterns of activity in a relatively more complex neural network as the target answer vectors. For example, local neural networks devices **2305** can be programmed to perform one or more iterative training techniques (e.g., gradient descent or stochastic gradient descent). In other implementations, the source approximators in local neural networks devices **2305** are trainable by, e.g., a dedicated training system or by a training system that is installed on a personal computer that can interact with the local neural networks devices **2305** to train source approximators.

(162) Each local neural networks device **2305** includes one or more wireless or wired data communication components. In the illustrated implementation, each local neural networks device **2305** includes at least one wireless data communication components such as a mobile phone transceiver, a wireless transceiver, or both. The mobile phone transceivers are able to exchange data with phone base station **2310**. The wireless transceivers are able to exchange data with a wireless access point **2315**. Each local neural networks device **2305** may also be able to exchange data with peer mobile computing devices.

(163) Phone base station **2310** and wireless access point **2315** are connected for data communication with one or more data communication networks **2325** and can exchange information with a server system **2320** over the network(s). Local neural networks devices **2305** are thus generally also in data communication with server system **2320**. However, this is not necessarily the case. For example, in implementations where local neural networks devices **2305** are trained by other data processing devices, local neural networks devices **2305** need only be in data communication with these other data processing devices at least once.

(164) Server system **2320** is a system of one or more data processing devices that is programmed to perform data processing activities in accordance with one or more sets of machine-readable instructions. The activities can include serving training sets to training systems for mobile computing devices **2305**. As discussed above, the training systems can be internal to mobile local neural networks devices **2305** themselves or on one or more other data processing devices. The training sets can include representations of the occurrence of topological structures corresponding to activity in a source neural network and corresponding input data.

(165) In some implementations, server system **2320** also includes the source neural network. However, this is not necessarily the case and server system **2320** may receive the training sets from yet another system of data processing device(s) that implement the source neural network.

(166) In operations, after server system **2320** receives a training set (from a source neural network that is found at server system **2320** itself or elsewhere), server system **2320** can serve the training set to trainers that train mobile computing devices **2305**. The source approximators in target local neural networks devices **2305** can be trained using the training set so that the target neural networks approximate the operations of the source neural network.

(167) FIGS. **24**, **25**, **26**, **27** are schematic illustrations of the use of representations of the occurrence of topological structures in the activity in a neural network in four different systems **2400**, **2500**, **2600**, **2700**. Systems **2400**, **2500**, **2600**, **2700** can be configured to perform any of a number of different operations. For example, systems **2400**, **2500**, **2600**, **2700** can perform object localization operations, object detection operations, object segmentation operations, object detection operations, prediction operations, action selection operations, or the like.

(168) Object localization operations locate an object within an image. For example, a bounding box can be constructed around an object. In some cases, object localization can be combined with object recognition, in which the localized object is labeled with an appropriate designation.

(169) Object detection operations classify image pixels as either belonging to a particular class (e.g., belonging to an object interest) or not. In general, object detection is performed by grouping pixels and forming bounding boxes around the pixel groups. The bounding box should be a tight fit around the object.

(170) Object segmentation generally assigns class labels to each image pixel. Thus, rather than a bounding box, object segmentation proceeds on a pixel-by-pixel basis and generally requires that only a single label be assigned to each pixel.

(171) Prediction operations seek to draw conclusions that are outside the range of a observed data. Although prediction operations can seek to forecast future occurrences (e.g., based on information about the past and current state), prediction operations can also seek to draw conclusions about the past and current state based on incomplete information on those states.

(172) Action selection operations seek to choose an action based on a set of conditions. Action selection operations have traditionally be broken down into different approaches such a symbol-based systems (classical planning), distributed solutions, and reactive or dynamic planning.

(173) Classification systems **2400**, **2500** each perform a desired operation on representations of the patterns of activity in a neural network. Systems **2600**, **2700** each perform a desired operation on approximations of representations of the patterns of activity in a neural network. In systems **2400**, **2500**, the patterns of activity that are represented occur in and are read from a source neural network device **1705** that is part of the system **2400**, **2500**. In contrast, in systems **2400**, **2500**, the patterns of activity that are approximately represented occur in a source neural network device that is not part of the system **2400**, **2500**. Nevertheless, the approximation of the representation of those patterns of activity are read from an approximator **1905** that is part of systems **2400**, **2500**.

(174) In additional detail, turning to FIG. **24**, system **2400** includes a source neural network **1705** and a linear processor **2410**. Linear processor **2410** is a device that performs operations based on a linear combination of the characteristics of representations of the patterns of activity in a neural network (or approximations of such representations). The operations can be, e.g., object localization operations, object detection operations, object segmentation operations, object detection operations, prediction operations, action selection operations, or the like.

(175) Linear processor **2410** includes an input **2420** and an output **2425**. Input **2420** is coupled to receive representations of the patterns of activity in source neural network **1705**. Linear processor **2410** can receive the representations of the patterns of activity in source neural network **1705** in a variety of ways. For example, the representations of the patterns of activity can be received as discrete events or as a continuous stream over a real time or non-real time communication channel. Output **2525** is coupled to output the processing result from linear processor **2410**. In some implementations, linear processor **2410** can be implemented on one or more computing devices with relatively limited computational performance. For example, linear processor **2410** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet.

(176) Turning to FIG. **24**, system **2400** includes a source neural network **1705** and a linear processor **2410**. Linear processor **2410** is a device that performs operations based on a linear combination of the characteristics of representations of the patterns of activity in a neural network (or approximations of such representations). The operations can be, e.g., object localization operations, object detection operations, object segmentation operations, prediction operations, action selection operations, or the like.

(177) Linear processor **2410** includes an input **2420** and an output **2425**. Input **2420** is coupled to receive representations of the patterns of activity in source neural network **1705**. Linear processor **2410** can receive the representations of the patterns of activity in source neural network **1705** in a variety of ways. For example, the representations of the patterns of activity can be received as

discrete events or as a continuous stream over a real time or non-real time communication channel. Output **2525** is coupled to output the processing result from linear processor **2410**. In some implementations, linear processor **2410** can be implemented on one or more computing devices with relatively limited computational performance. For example, linear processor **2410** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet. (178) In FIG. **25**, classification system **2500** includes source neural network **1705** and a neural network **2510**. Neural network **2510** is a neural network device that is configured to perform operations based on a non-linear combination of the characteristics of representations of the patterns of activity in a neural network (or approximations of such representations). The operations can be, e.g., object localization operations, object detection operations, object segmentation operations, prediction operations, action selection operations, or the like. In the illustrated implementation, neural network **2510** is a feedforward network that includes an input layer **2520** and an output layer **2525**. As with linear processor **2410**, neural network **2510** can receive the representations of the patterns of activity in source neural network **1705** in a variety of ways.

(179) In some implementations, neural network **2510** can perform inferences on one or more computing devices with relatively limited computational performance. For example, neural network **2510** can be implemented on a personal computer or a mobile computing device such as a smart phone or tablet, e.g., in a Neural Processing Unit of such a device. Like system **2400**, system **2500** will generally be a dispersed system in which a remote neural network **2510** communicates with source neural network **1705**, e.g., via a data communications network. In some implementations, neural network **2510** can be, e.g., a deep neural network such as a convolutional neural network.

(180) In FIG. **26**, system **2600** includes source approximator **1905** and a linear processor **2410**. Despite any differences between approximation **1200'** representation **1200**, processor **2410** can still perform operations on approximation **1200'**.

(181) In FIG. **27**, system **2700** includes source approximator **1905** and a neural network **2510**. Despite any differences between approximation **1200'** and representation **1200**, neural network **2510** can still perform operations on approximation **1200'**.

(182) In some implementations, systems **2600**, **2700** can be implemented on an edge device, such as, e.g., edge devices **2100**, **2200** (FIGS. **21**, **22**). In some implementations, systems **2600**, **2700** can be implemented as part of a system in which local neural networks can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network, such as system **2300** (FIG. **23**).

(183) FIG. **28** is a schematic illustration of a reinforcement learning system **2800** that includes an artificial neural network that can be trained using representations of the occurrence of topological structures corresponding to activity in a source neural network. Reinforcement learning is a type of machine learning in which an artificial neural network learns from feedback regarding the consequences of actions taken in response to the artificial neural network's decisions. A reinforcement learning system moves from one state in the environment to another by performing actions and receiving information characterizing the new state and a reward and/or regret that characterizes the success (or lack of success) of the action. Reinforcement learning seek to maximize the total reward (or minimize the regret) through a learning process.

(184) In the illustrated implementation, the artificial neural network in reinforcement learning system **2800** is a deep neural network **2805** (or other deep learning architecture) that is trained using a reinforcement learning approach. In some implementations, deep neural network **2805** can be a local artificial neural network (such as neural network **2510** (FIGS. **25**, **27**), and implemented locally on, e.g., an automobile, a plane, a robot, or other device. However, this is not necessarily the case and in other implementations, deep neural network **2805** can be implemented on a system of networked devices.

(185) In addition to a source approximator **1905** and deep neural network **2805**, reinforcement learning system **2800** includes an actuator **2810**, one or more sensors **2815**, and a teacher module

2820. In some implementations, reinforcement learning system **2800** also includes one or more sources **2825** of additional data.

(186) Actuator **2810** is a device that controls a mechanism or a system that interacts with an environment **2830**. In some implementations, actuator **2810** controls a physical mechanism or system (e.g., the steering of an automobile or the positioning of a robot). In other implementations, actuator **2810** can control a virtual mechanism or system (e.g., a virtual game board or an investment portfolio). Thus, environment **2830** may also be physical or virtual.

(187) Sensor(s) **2815** are devices that measure characteristics of the environment **2830**. At least some of the measurements with characterize interactions between the controlled mechanism or system and other aspects of the environment **2830**. For example, when actuator **2810** steers an automobile, sensor(s) **2815** may measure one or more of the speed, direction, and acceleration of the automobile, the proximity of the automobile to other features, and the response of other features to the automobile. As another example, when actuator **2810** controls an investment portfolio, sensor(s) **2815** may measure the value and risk associated with the portfolio.

(188) In general, both source approximator **1905** and teacher module **2820** are coupled to receive at least some of the measurements made by sensor(s) **2815**. For example, source approximator **1905** can receive measurement data at input layer **1915** and output an approximation **1200'** of a representation of topological structures that arise in the patterns of activity in a source neural network.

(189) Teacher module **2820** is a device that is configured to interpret the measurements received from sensor(s) **2815** and provide a reward and/or a regret to deep neural network **2805**. Rewards are positive and indicate successful control of the mechanism or system. Regrets are negative and indicate unsuccessful or less than optimal control. In general, teacher module **2820** also provides a characterization of the measurements along with the reward/regret for reinforcement learning. In general, the characterization of the measurements is an approximation of a representation of topological structures that arise in the patterns of activity in a source neural network (such as approximation **1200'**). For example, teacher module **2820** may read approximations **1200'** output from source approximator **1905** and pair the read approximations **1200'** with corresponding reward/regret values.

(190) In many implementations, reinforcement learning does not occur in system **2800** in real time or during active control of actuator **2810** by deep neural network **2805**. Rather, training feedback can be collected by teacher module **2820** and used for reinforcement training when deep neural network **2805** is not actively instructing actuator **2810**. For example, in some implementations, teacher module **2820** can be remote from deep neural network **2805** and only in intermittent data communication with deep neural network **2805**. Regardless of whether reinforcement learning is intermittent or continuous, deep neural network **2805** can be evolved, e.g., to optimize reward and/or reduce regret using the information received from teacher module **2820**.

(191) In some implementations, system **2800** also includes one or more sources **2825** of additional data. Source approximator **1905** can also receive data from data sources **2825** at input layer **1915**. In these instances, approximation **1200'** will result from processing both sensor data and the data from data sources **2825**.

(192) In some implementations, the data collected by one reinforcement learning system **2800** can be used for training or reinforcement learning of other systems, including other reinforcement learning systems. For example, the characterization of the measurements along with the reward/regret values can be provided by teacher module **2820** to a data exchange system that collects such data from a variety of reinforcement learning systems and redistributes it among them. Further, as discussed above, the characterization of the measurements can be an approximation of a representation of topological structures that arise in the patterns of activity in a source neural network, such as approximation **1200'**.

(193) The particular operations that are performed by reinforcement learning system **2800** will of

course depend on the particular operational context. For example, in contexts where source approximator **1905**, deep neural network **2805**, actuator **2810**, and sensors **2815** are part of an automobile, deep neural network **2805** can perform object localization and/or detection operations while steering the automobile.

(194) In implementations where the data collected by reinforcement learning system **2800** is used for training or reinforcement learning of other systems, reward/regret values and approximations **1200'** that characterize the state of the environment when object localizations and/or detection operations were performed can be provided to the data exchange system. The data exchange system can then distribute the reward/regret values and approximations **1200'** to other reinforcement learning systems **2800** associated with other vehicles for reinforcement learning at those other vehicles. For example, reinforcement learning can be used to improve object localization and/or detection operations at a second vehicle using the reward/regret values and approximations **1200'**.

(195) However, the operations that are learned at other vehicles need not be identical to the operations that are performed by deep neural network **2805**. For example, reward/regret values that are based on travel time and approximations **1200'** that result from the input of sensor data characterizing, e.g., an unexpectedly wet road at a location identified by a GPS data source **2825** can be used for route planning operations at another vehicle.

(196) Embodiments of the subject matter and the operations described in this specification can be implemented in digital electronic circuitry, or in computer software, firmware, or hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them. Embodiments of the subject matter described in this specification can be implemented as one or more computer programs, i.e., one or more modules of computer program instructions, encoded on computer storage medium for execution by, or to control the operation of, data processing apparatus. Alternatively or in addition, the program instructions can be encoded on an artificially-generated propagated signal, e.g., a machine-generated electrical, optical, or electromagnetic signal, that is generated to encode information for transmission to suitable receiver apparatus for execution by a data processing apparatus. A computer storage medium can be, or be included in, a computer-readable storage device, a computer-readable storage substrate, a random or serial access memory array or device, or a combination of one or more of them. Moreover, while a computer storage medium is not a propagated signal, a computer storage medium can be a source or destination of computer program instructions encoded in an artificially-generated propagated signal. The computer storage medium can also be, or be included in, one or more separate physical components or media (e.g., multiple CDs, disks, or other storage devices).

(197) The operations described in this specification can be implemented as operations performed by a data processing apparatus on data stored on one or more computer-readable storage devices or received from other sources.

(198) The term “data processing apparatus” encompasses all kinds of apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, a system on a chip, or multiple ones, or combinations, of the foregoing. The apparatus can include special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application-specific integrated circuit). The apparatus can also include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, a cross-platform runtime environment, a virtual machine, or a combination of one or more of them. The apparatus and execution environment can realize various different computing model infrastructures, such as web services, distributed computing and grid computing infrastructures.

(199) A computer program (also known as a program, software, software application, script, or code) can be written in any form of programming language, including compiled or interpreted languages, declarative or procedural languages, and it can be deployed in any form, including as a

stand-alone program or as a module, component, subroutine, object, or other unit suitable for use in a computing environment. A computer program may, but need not, correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub-programs, or portions of code). A computer program can be deployed to be executed on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

(200) The processes and logic flows described in this specification can be performed by one or more programmable processors executing one or more computer programs to perform actions by operating on input data and generating output. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application-specific integrated circuit).

(201) Processors suitable for the execution of a computer program include, by way of example, both general and special purpose microprocessors, and any one or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read-only memory or a random access memory or both. The essential elements of a computer are a processor for performing actions in accordance with instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto-optical disks, or optical disks. However, a computer need not have such devices. Moreover, a computer can be embedded in another device, e.g., a mobile telephone, a personal digital assistant (PDA), a mobile audio or video player, a game console, a Global Positioning System (GPS) receiver, or a portable storage device (e.g., a universal serial bus (USB) flash drive), to name just a few. Devices suitable for storing computer program instructions and data include all forms of non-volatile memory, media and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto-optical disks; and CD-ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(202) To provide for interaction with a user, embodiments of the subject matter described in this specification can be implemented on a computer having a display device, e.g., a CRT (cathode ray tube) or LCD (liquid crystal display) monitor, for displaying information to the user and a keyboard and a pointing device, e.g., a mouse or a trackball, by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well; for example, feedback provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback; and input from the user can be received in any form, including acoustic, speech, or tactile input. In addition, a computer can interact with a user by sending documents to and receiving documents from a device that is used by the user; for example, by sending web pages to a web browser on a user's client device in response to requests received from the web browser.

(203) While this specification contains many specific implementation details, these should not be construed as limitations on the scope of any inventions or of what may be claimed, but rather as descriptions of features specific to particular embodiments of particular inventions. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable subcombination. Moreover, although features may be described above as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed

combination may be directed to a subcombination or variation of a subcombination.

(204) Similarly, while operations are depicted in the drawings in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the embodiments described above should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

(205) Thus, particular implementations of the subject matter have been described. Other implementations are within the scope of the following claims. In some cases, the actions recited in the claims can be performed in a different order and still achieve desirable results. In addition, the processes depicted in the accompanying figures do not necessarily require the particular order shown, or sequential order, to achieve desirable results. In certain implementations, multitasking and parallel processing may be advantageous.

(206) A number of embodiments have been described. Nevertheless, it will be understood that various modifications may be made. For example, although representation **1200** is a binary representation in which each bit individually represents the presence or absence of a feature in a graph, other representations of information are possible. For example, a vector or matrix of multi-valued, non-binary digits can be used to represent, e.g., the presence or absence of features and possibly other characteristics of those features. An example of such a characteristic is the weight of edges with the activity that constitutes the features.

(207) Accordingly, other embodiments are within the scope of the following claims.

Claims

1. A computer-implemented method comprising: training a functional state of a recurrent artificial neural network that includes two thousand nodes or more to propagate an input according to a learned signal flow within the recurrent artificial neural network until processing of the input reaches a predefined end and the recurrent artificial neural network returns to a state in which only background or no signal transmission activity occurs, the training comprising: structuring the recurrent artificial neural network to obtain a trained recurrent artificial neural network, the structuring comprising: reading a plurality of digits output from the recurrent artificial neural network, wherein each digit represents whether or not activity within a particular group of nodes in the recurrent artificial neural network comports with a respective pattern of activity; and evolving a structure of the recurrent artificial neural network, wherein evolving the structure of the recurrent artificial neural network comprises: characterizing a complexity of each pattern of signal transmission activity in the structure, wherein the complexity is measured based on simplex counts or Betti numbers of the pattern, using the characterization as an indication of whether the structure has optimized a reward, and iteratively evolving the structure if the reward has not been optimized; processing the input by the trained recurrent artificial neural network, the input comprising a stream of video or audio data; monitoring complexities of patterns of signal transmission activity that occur between three or more nodes in the trained recurrent artificial neural network in response to the processing of the input, the monitoring comprising, during the processing of the input: outputting, by the trained recurrent artificial neural network, a second plurality of digits, wherein a value of each digit in the second plurality of digits represents whether or not signal transmission activity that occurs between the three or more nodes in the trained recurrent artificial neural network in response to the processing of the input comports with a respective pattern of signal transmission activity, regardless of a location at which the respective pattern of signal transmission activity occurs in the trained recurrent artificial neural network, wherein the second plurality of digits enables an

isomorphic topological reconstruction of a graph of the trained recurrent artificial neural network; determining the complexity of each pattern of signal transmission activity; identifying a time at which an object or sound is recognized in the stream of video or audio data, comprising identifying, based on the complexity of each pattern of signal transmission activity, a distinguishable level of complexity of the signal transmission activity, wherein the distinguishable level of complexity comprises an upward deviation, downward deviation, or change of complexity from a predetermined threshold level of complexity; outputting, at the identified time and before the trained recurrent artificial neural network completes the processing of the input to reach the predefined end and return to the state in which only background or no signal transmission activity occurs, the patterns of signal transmission activity that have the distinguishable level of complexity; and transmitting, encrypting, or storing the output patterns of signal transmission activity that have the distinguishable level of complexity; and allowing the trained recurrent artificial neural network to return to the state in which only background or no signal transmission activity occurs.

2. The method of claim 1, further comprising determining whether or not each group of nodes comports with the respective pattern of signal transmission activity by identifying clique patterns of signal transmission activity of the trained recurrent artificial neural network.

3. The method of claim 2, wherein the method further comprises defining a plurality of windows of time during which the signal transmission activity of the trained recurrent artificial neural network is responsive to the input into the trained recurrent artificial neural network, wherein each digit represents whether or not the signal transmission activity within the group of nodes comports with a respective clique pattern of signal transmission activity during a respective window of the plurality of windows.

4. The method of claim 3, wherein the method further comprises identifying a first window of time within the plurality of windows of time based on an increased likelihood of the clique patterns being identified during the first window of time when compared to other windows of the plurality of windows of time.

5. The method of claim 1, wherein each of the respective patterns of signal transmission activity encloses a cavity.

6. The method of claim 2, wherein identifying cliques comprises discarding or ignoring cliques that are present in other cliques.

7. The method of claim 1, wherein each of the digits is output from a reader node coupled to a collection of nodes in the trained recurrent artificial neural network, wherein the reader nodes indicate if signal transmission activity of the nodes in the collection comports with a particular pattern of signal transmission activity.

8. The method of claim 1, further comprising transmitting or storing a quantity of the digits, wherein transmitting or storing a quantity of the digits comprises: transmitting or storing digits associated with patterns of signal transmission activity having a complexity higher than the predetermined threshold level; and discarding or ignoring digits associated with patterns of signal transmission activity having a complexity lower than the predetermined threshold level.

9. The method of claim 1, wherein times at which the patterns of signal transmission activity that have the distinguishable level of complexity occur provide a classification as to whether data in the data stream at the times satisfies object or sound recognition criteria.

10. A data transmitter or data storage device comprising an encoder comprising one or more computers operable to perform operations, the operations comprising: training a functional state of a recurrent artificial neural network that includes two thousand nodes or more to propagate an input according to a learned signal flow within the recurrent artificial neural network until processing of the input reaches a predefined end and the recurrent artificial neural network returns to a state in which only background or no signal transmission activity occurs, the training comprising: structuring the recurrent artificial neural network to obtain a trained recurrent artificial neural network, the structuring comprising: reading a plurality of digits output from the recurrent artificial

neural network, wherein each digit represents whether or not activity within a particular group of nodes in the recurrent artificial neural network comports with a respective pattern of activity; and evolving a structure of the recurrent artificial neural network, wherein evolving the structure of the recurrent artificial neural network comprises: characterizing a complexity of each pattern of signal transmission activity in the structure, wherein the complexity is measured based on simplex counts or Betti numbers of the pattern, using the characterization as an indication of whether the structure has optimized a reward, and iteratively evolving the structure if the reward has not been optimized; encoding a result of processing of an input by the trained recurrent artificial neural network, the input comprising a stream of video or audio data, wherein encoding the result of processing comprises: processing the input by the trained recurrent artificial neural network; monitoring complexities of patterns of signal transmission activity that occur between three or more nodes in the trained recurrent artificial neural network in response to the processing of the input, the monitoring comprising, during the processing of the input: outputting, by the trained recurrent artificial neural network, a second plurality of digits from the trained recurrent artificial neural network, wherein a value of each digit in the second plurality of digits represents whether or not signal transmission activity that occurs between the three or more nodes in the trained recurrent artificial neural network in response to the processing of the input comports with a respective pattern of signal transmission activity, regardless of a location at which the respective pattern of signal transmission activity occurs in the trained recurrent artificial neural network, wherein the second plurality of digits enables an isomorphic topological reconstruction of a graph of the trained recurrent artificial neural network, determining the complexity of each pattern of signal transmission activity; identifying a time at which an object or sound is recognized in the stream of video or audio data, comprising identifying, based on the complexity of each pattern of signal transmission activity, a distinguishable level of complexity of the signal transmission activity, wherein the distinguishable level of signal transmission activity comprises an upward deviation, downward deviation, or change of complexity from a predetermined threshold level of complexity; outputting, at the identified time and before the trained recurrent artificial neural network completes the processing of the input to reach the predefined end and return to the state in which only background or no signal transmission activity occurs, the result of processing of the input, the result comprising the patterns of signal transmission activity that have the distinguishable level of complexity; and transmitting, encrypting, or storing the output patterns of signal transmission activity that have the distinguishable level of complexity; and allowing the trained recurrent artificial neural network to return to the state in which only background or no signal transmission activity occurs.

11. The encoder of claim 10, wherein the operations further comprise determining whether or not each group of nodes comports with the respective pattern of signal transmission activity by identifying clique patterns of signal transmission activity of the trained recurrent artificial neural network.

12. The encoder of claim 11, wherein the operations further comprise defining a plurality of windows of time during which the signal transmission activity of the trained recurrent artificial neural network is responsive to the input into the trained recurrent artificial neural network, wherein each digit represents whether or not signal transmission activity within the group of nodes comports with a respective clique pattern of signal transmission activity during a respective window of the plurality of windows.

13. The encoder of claim 12, wherein the operations further comprise identifying a first window of time within the plurality of windows of time based on an increased likelihood of the clique patterns being identified during the first window of time when compared to other windows of the plurality of windows of time.

14. The encoder of claim 10, wherein the respective patterns of signal transmission activity enclose cavities.

15. The encoder of claim 10, wherein each of the digits is output from a reader node coupled to a collection of nodes in the trained recurrent artificial neural network, wherein the reader nodes indicate if signal transmission activity of the nodes in the collection comports with a particular pattern of signal transmission activity.

16. The encoder of claim 10, further comprising transmitting or storing a quantity of the digits, wherein transmitting or storing a quantity of the digits comprises: transmitting or storing digits associated with patterns of signal transmission activity having a complexity higher than the predetermined threshold level; and discarding or ignoring digits associated with patterns of signal transmission activity having a complexity lower than the predetermined threshold level.

17. The encoder of claim 10, wherein the encoder is a video encoder.

18. The encoder of claim 10, wherein times at which the patterns of signal transmission activity that have the distinguishable level of complexity occur provide a classification as to whether data in the data stream at the times satisfies object or sound recognition criteria.
