

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12412231
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Elzur; Uri

Graphics processing unit with network interfaces

Abstract

Examples described herein relate to a first graphics processing unit (GPU) comprising at least one integrated network interface controller (NIC), wherein a data plane of the at least one integrated NIC is configured by a particular entity. In some examples, the data plane of the at least one integrated NIC is configured by the particular entity as a control plane but the first GPU and a central processing unit (CPU) are prevented from configuring the data plane of the at least one integrated NIC.

Inventors:	Elzur; Uri (San Jose, CA)
Applicant:	Intel Corporation (Santa Clara, CA)
Family ID:	83602495
Assignee:	Intel Corporation (Santa Clara, CA)
Appl. No.:	17/853793
Filed:	June 29, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20220335563 A1	Oct. 20, 2022

Related U.S. Application Data

us-provisional-application US 63218832 20210706

Publication Classification

Int. Cl.: G06T1/20 (20060101); G06F13/42 (20060101); G06T1/60 (20060101); H04L41/08 (20220101); H04L41/0803 (20220101)

U.S. Cl.:

CPC **G06T1/20** (20130101); **G06F13/4282** (20130101); **G06T1/60** (20130101);
H04L41/0803 (20130101); G06F2213/3808 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6393483	12/2001	Latif	370/359	H04L 41/22
9600429	12/2016	Pope	N/A	G06F 13/4282
9602437	12/2016	Bernath	N/A	H04L 49/9084
10997106	12/2020	Bandaru	N/A	G06F 13/385
11150963	12/2020	Nainar	N/A	G06F 9/541
11531752	12/2021	Abodunrin	N/A	G06F 21/53
11645104	12/2022	Rajagopal	718/1	G06F 9/45558
11863376	12/2023	Ang	N/A	H04L 41/0803
11995024	12/2023	Ang	N/A	G06F 13/387
12206554	12/2024	Choudhury	N/A	H04L 41/0893
2007/0139423	12/2006	Kong et al.	N/A	N/A
2007/0273699	12/2006	Sasaki et al.	N/A	N/A
2010/0013839	12/2009	Rawson	345/520	G06F 3/14
2014/0098113	12/2013	Sabony	345/519	G06T 1/20
2014/0223131	12/2013	Agarwal et al.	N/A	N/A
2016/0337193	12/2015	Rao	N/A	H04L 67/63
2017/0180272	12/2016	Bernath	N/A	H04L 49/901
2018/0314638	12/2017	LeBeane	N/A	G06F 12/1081
2018/0349062	12/2017	Sines	N/A	N/A
2019/0042741	12/2018	Abodunrin	N/A	G06F 13/4068
2020/0226096	12/2019	Koker et al.	N/A	N/A
2020/0278892	12/2019	Nainar	N/A	H04L 67/10
2021/0203704	12/2020	Pohl	N/A	A63F 13/355
2022/0085916	12/2021	Debbage et al.	N/A	N/A

2022/0100491	12/2021	Voltz	N/A	G06F 8/65
2022/0197681	12/2021	Rajagopal	N/A	G06F 9/45558
2023/0041806	12/2022	Singh	N/A	H04L 63/0428
2023/0195675	12/2022	Ang	710/313	G06F 13/387
2023/0198833	12/2022	Ang	709/220	H04L 41/0806

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
102413031	12/2011	CN	N/A
WO-2020172692	12/2019	WO	G06F 9/5077
2020190798	12/2019	WO	N/A
2020190801	12/2019	WO	N/A

OTHER PUBLICATIONS

‘UNO: Unifying Host and Smart NIC Offload for Flexible Packet Processing’ by Yanfang Le et al., SoCC '17, Sep. 24-27, 2017. (Year: 2017). cited by examiner

‘P4NFV: P4 Enabled NFV Systems with SmartNICs’ by Ali Mohammadkhan et al., 2019 IEEE Conference on Network Function Virtualization and Software Defined Networks (NFV-SDN). (Year: 2019). cited by examiner

‘HyperV: A High Performance Hypervisor for Virtualization of the Programmable Data Plane’ by Cheng Zhang et al., 2017 26th International Conference on Computer Communication and Networks (ICCCN), Jul. 2017. (Year: 2017). cited by examiner

‘HyperVDP: High-Performance Virtualization of the Programmable Data Plane’ by Cheng Zhang et al., IEEE Journal On Selected Areas in Communications, vol. 37, No. 3, Mar. 2019. (Year: 2019). cited by examiner

‘Enhancing Networks with SmartNICs’ by Mellanox, copyright 2018. (Year: 2018). cited by examiner

AWS, “Amazon EC2 P4 Instances”, <https://aws.amazon.com/ec2/instance-types/p4/>, Nov. 2, 2020, 8 pages. cited by applicant

Bouffler, Brendan, “In the search for performance, there's more than one way to build a network”, AWS HPC Blog, <https://aws.amazon.com/blogs/hpc/in-the-search-for-performance-theres-more-than-one-way-to-build-a-network/>, Jun. 22, 2021, 5 pages. cited by applicant

Gebara, Nadeen et al., “In-Network Aggregation for Shared Machine Learning Clusters”, Proceedings of the 4th MLSys Conference, San Jose, CA, USA, 2021, 16 pages. cited by applicant
NVIDIA, “DGX Station A100 User Guide”, DU-10189-001 v5.0.2, <https://docs.nvidia.com/dgx/dgx-station-a100-user-guide/index.html>, Nov. 2022, 74 pages. cited by applicant

Singhvi, Arjun et al., “1RMA: Re-envisioning Remote Memory Access for Multi-tenant Datacenters”, SIGCOMM '20, Aug. 10-14, 2020, Virtual Event, NY, USA, 14 pages. cited by applicant

International Search Report and Written Opinion for PCT Patent Application No.

PCT/US22/36397, Mailed Oct. 24, 2022, 12 pages. cited by applicant

U.S. Appl. No. 17/853,711, filed Jun. 29, 2022, 71 pages. cited by applicant

Notice of Allowance for U.S. Appl. No. 17/853,711, Mailed Aug. 29, 2024, 10 pages. cited by

applicant

Klenk M Benjamin: "Communication Architectures for Scalable GPU-centric Computing Systems", Dec. 31, 2017 (Dec. 31, 2017), XP055774858, Retrieved from the Internet:

URL:https://archiv.ub.uni-heidelberg.de/volltextserver/23954/1/bklenk_dissertation_final.pdf

[retrieved on Feb. 11, 2021], 217 pages. cited by applicant

Partial European Search Report for Patent Application No. 22838425.1, Mailed Apr. 24, 2025, 14 pages. cited by applicant

Primary Examiner: Snyder; Steven G

Attorney, Agent or Firm: Compass IP Law PC

Background/Summary

RELATED APPLICATION (1) This application claims priority to U.S. Provisional Application No. 63/218,832, filed Jul. 6, 2021. The entire contents of that application are incorporated by reference in its entirety.

BACKGROUND

(1) In Infrastructure-as-a-service (IaaS) or as part of a composable Data Center (DC), applications and other software share utilization of central processing units (CPUs), field programmable gate arrays (FPGAs), graphics processing units (GPUs), memory devices, and other devices. As Deep Learning (DL) training workloads grow in size, even exponentially, multiple GPUs are utilized. The size, scale, and complexity of modern DL workloads training leads to use of a dedicated scale out fabric (e.g., an internodal connectivity scheme) to provide network bandwidth for DL training workloads. For sharing an installed GPU cluster with multiple tenants and/or for workloads shared across multiple GPUs, some GPUs are attached to a Cloud Service Provider (CSP) DC network and use smartNICs (e.g., Amazon Web Service (AWS) P4 d instance) to access the DC network, in compliance with CSP access policies. Other CSPs utilize a proprietary pod with a technology different than the DC network technology with higher bandwidth for training workloads.

(2) Current approaches to connect a GPU server to a CSP DC network include utilization of scale out network interface controllers (NICs) and ability to support the varying requirements of different CSPs who may deploy different technologies, different topologies and/or different security/privacy/isolation policies. As the network space evolves rapidly, the requirements and technologies may frequently change. However, many of these approaches utilize smartNICs, that can incur relatively high cost, high power, and high board space, but can deliver relatively small bandwidth for GPU-to-GPU communications and for GPU-to-CPU communications over the DC network.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 depicts an example system with GPU and integrated NICs.

(2) FIG. 2 depicts an example system.

(3) FIG. 3 depicts an example system.

(4) FIG. 4 depicts an example system.

(5) FIG. 5 depicts an example of a GPU server.

(6) FIG. 6 depicts an example process.

(7) FIGS. 7A-7D depict an example GPU.

(8) FIG. 8 depicts an example network interface.

(9) FIG. 9 depicts a system.

(10) FIG. 10 depicts a system.

DETAILED DESCRIPTION

(11) At least to provide for communication among GPUs, or other processors or devices (e.g., at least accelerators, FPGAs, application specific integrated circuits (ASICs), CPUs, and so forth), one or more NICs can be integrated into a GPU so that the same semiconductor device or system on chip (SoC) can include a GPU and one or more NICs. A configured smartNIC that communicates in compliance with a CSP's policies for DC network access can be coupled to NICs that are integrated into the GPU or is directly coupled (or assigned to the CPU while connected over an interface or another network device). For example, a GPU NIC can refer to or include at least one NIC that is integrated into a GPU. Note that reference to a CSP can refer in addition, or alternatively to a communication service provider (CoSP). The CSP configured smart NIC can configure one or more NICs integrated into the GPU. For example, a data plane of the one or more integrated NICs can be configured by a CSP controlled smartNIC in one or more of: reliability protocol to apply, packet encryption to apply, packet decryption to apply, network protocol or frame format to utilize, and/or permitted destination for packets.

(12) One or more NICs integrated into the GPU can provide reduced power and performance for high bandwidth and low latency workloads. Capabilities of an integrated NIC can be augmented by a bump-in-the-wire processor accelerator, and/or smartNIC to provide additional features for a CSP or CoSP. One or more integrated NICs may communicate with other NICs based on proprietary protocols or standardized remote direct memory access (RDMA) protocols (e.g., InfiniBand, remote direct memory access (RDMA) over Converged Ethernet (RoCE), RoCEv2, remote memory access (RMA), a specialized proprietary subset of RDMA for its data transfers, and so forth). For example, one or more integrated NICs can be used to provide communication bandwidth for artificial intelligence (AI) training or inference traffic with higher bandwidth, lower latency and other parameters independent from other DC traffic. A GPU can utilize a GPU NIC, NIC coupled through a device interface, or indirectly coupled assigned NIC to communication with another GPU or device. The GPU can access a memory through a NIC via a network.

(13) A CSP or other entity can provide a control plane for a plurality of data planes of one or more NICs of a GPU. One or more smartNICs can be deployed for CSPs to configure one or more GPU NIC data planes. In some examples, the one or more NICs of a GPU provides isolated control plane access to an external CoSP or CSP NIC or smartNIC or other designated management entity so that access can be authenticated to ensure only the appropriate designated management entity can configure and control the GPU NICs. In some examples, the GPU and the local CPU are prevented from accessing, altering, or editing the configuration of the integrated NIC into the GPU or NIC directly attached to the GPU. A CSP programmed control plane can configure data plane(s) of NIC(s) integrated inside the GPU or inside another device (e.g., a switch on a GPU node). An external trusted configuration entity can configure control and configure the CSP programmed control plane. The external trusted configuration entity can be an entity designated by a CSP or CoSP. The external trusted configuration entity can include a software defined networking (SDN) controller, orchestrator, administrator, BMC, container, virtual machine (VM), microservice resident software, or others.

(14) CSP specific control planes can be implemented in one or more trusted network interface devices or smartNICs. For example, for a CSP, a smartNIC can act as a trusted entity and configure a control plane of one or more integrated NICs of the GPU. A CSP controlled control plane device can provide direct link for exclusive programmability of one or more of the GPU integrated data planes (e.g., out-of-band) or by a link shared with data transmission as well (e.g., in-band). In some examples, one or more integrated NICs of the GPU are isolated from the server operating system

(OS), hypervisor tenant/user or application and/or from other integrated NICs of the GPU, so that control plane configurations of NICs are separately controlled, thereby providing the CSP with sole capability to control the GPU NICs.

(15) Different CoSPs or CSPs can utilize their own protocols, policies (e.g., congestion, reliability) and can deploy preferred policies. SmartNIC processor core(s) can execute one or more CoSP's or CSP's management software and implement the one or more CoSP's or CSP's security policies. A CoSP or CSP programmed control plane (e.g., software defined networking (SDN)) interacting with a smartNIC can control the integrated data plane NIC (through a dedicated, secure link or virtual link) for network traffic scheduling, injection rate controls, network traffic congestion control, packet receipt reliability, security, memory isolation, etc. CoSP or CSP programmed control plane can participate-in and utilize CoSP or CSP preferred server/network tools (e.g., for inventory, provision, configuration, updates, routing, resource or NICs or tenant workload load balancing, management, telemetry gathering, failure recovery etc.), and/or latency and latency tail control measures.

(16) FIG. 1 depicts an example system with GPU and integrated NICs. In some examples, GPU compute **102** as well as memory **103** can be communicatively coupled to one or more NICs **110**. For example, GPU compute **102**, memory **103**, and NICs **110** can be communicatively coupled using an integrated or externally packaged switch and/or via interface **104** (e.g., Peripheral Component Interconnect express (PCIe), Compute Express Link (CXL), serializer de-serializer (SerDes), silicon photonics or optical interface, or other protocol or connectivity technology). However, an integrated or externally packaged switch can be optional. GPU compute **102**, memory **103**, and NICs **110** can be formed on a same integrated circuit chip or different integrated circuit chips, same die or different dies, or same package or different packages. GPU compute **102**, memory **103**, and NICs **110** can be communicatively coupled using chip-to-chip communications, die-to-die communications, packet-based communications, communications over a device interface, fabric-based communications, and so forth. Die-to-die communications can be consistent with Embedded Multi-Die Interconnect Bridge (EMIB), interposer, or other interfaces (e.g., Universal Chiplet Interconnect Express (UCIe), described at least in UCIe 1.0 Specification (2022), as well as earlier versions, later versions, and variations thereof).

(17) GPU compute **102** can include single-instruction multiple thread (SIMT) execution circuitry, one or more tensor cores, a thread dispatcher, or a thread scheduler. GPU compute **102** can include other circuitry and technologies described with respect to FIG. 7A, 7B, 7C, or 7D. In some examples, GPU compute **102** can be virtualized or composable as part of a composite node. Note that reference to GPU can instead refer to an XPU or an accelerator. GPU compute **102** can include or access memory **103** to read or write data. In some examples, memory **103** is in a separate device than that of GPU compute **102**.

(18) One or more instances of GPU compute **102** can be disaggregated and located on a same or different servers and provide a composite or composed GPU compute capabilities from the one or more instances of GPU compute via communications technologies such as fabrics, networks, interconnects, device interfaces, and so forth.

(19) NICs **110** can include multiple NICs **0** to **7**, although other numbers of NICs can be used. A NIC can include circuitries and technologies described with respect to FIG. 8. For example, a NIC can include direct memory access circuitry and direct memory access (DMA), congestion management (CM), multi-pathing (MP), reliability layer (RL), media access control and physical layer processing (MAC/PHY).

(20) FIG. 2 depicts an example system. GPU compute **102** (or a portion thereof) can utilize one or more NICs of NICs **208** to communicate with other devices. NICs **208** can be integrated into silicon of GPU compute **102**, in a same GPU package, or in a connectivity device such a switch, potentially serving one or more GPUs or a portion of GPU compute **102**. In some examples, optionally, switch **204** can utilized to provide communications between GPU compute **102** or

memory **103** and one or more NIC via connection **202** and interface **206**. Interface **206** can utilize one or more of the following standards and technology: Peripheral Component Interconnect express (PCIe), Compute Express Link (CXL), serializer de-serializer (SerDes), silicon photonics or optical interface, or other protocol or connectivity technology. GPU compute **102**, connection **202**, and NICs **208** can be formed on a same integrated circuit chip or different integrated circuit chips, same die or different dies, or same package or different packages.

(21) In some examples, data planes of one or more of NICs **0-7** can be configured by one or more different CoSPs or CSPs. CSP control plane **200** can provide configuration of one or more CSP data planes of one or more of NICs **0-7** via network interface device **210**. In some examples, network interface device **210** can include one or more of: a network interface controller (NIC), a remote direct memory access (RDMA)-enabled NIC, SmartNIC, router, switch, forwarding element, infrastructure processing unit (IPU), or data processing unit (DPU). In some examples, network interface device **210** and NICs **208** can communicate out-of-band over designated management interfaces (e.g., PCIe, baseboard management controller (BMC), over Distributed Management Task Force (DMTF) protocols, etc.). In some examples, network interface device **210** and NICs **208** can communicate in-band over one or more links **212** used for data transmission.

(22) CSP control plane **200** can configure network interface device **210** to configure data planes of one or more of NICs among NICs **208** in one or more of: reliability layer, security (e.g., encryption/decryption schemes to apply), network protocol and/or frame format options to apply. Control of the data plane can be exclusive and may not, in some examples, be changed by GPU compute **102**, software executed by GPU **102**, or CPU executed host software. Data plane programming can include selection and/or configuration of one or more of: congestion management (CM) to apply (e.g., Swift, Data Center Quantized Congestion Notification (DCQCN), Explicit Congestion Notification (ECN), Priority-based Flow Control (PFC), Data Center TCP (DCTCP), Timely, and so forth), or allow for one or more CSP-specific CM algorithms. Data plane programming can also allow for programming of the reliability layer (RL) to provide transport services (e.g., guaranteed packet delivery, in-order packet delivery, packet drop identification and/or packet duplication avoidance) and packet transport and packet re-transmission similar to services provided by TCP, UDP, or other reliability protocols. Data plane programming can also allow for programming of the multi pathing protocol (MP) to select a manner of utilizing available multiple routes or paths between sender and a receiver. Data plane programming can also allow for programming of a protocol to utilize for packet transmissions or receipt (e.g., RoCE, Scalable Reliable Datagram (SRD), QUIC), and so forth. For example, CSP control plane **200** can be implemented as an SDN controller that is executed on a same or different server than that of GPU compute **102**. Connectivity between CSP control plane **200** and switch **204** can occur based on packet based communications in accordance with standardized or proprietary protocols. For example, the SDN controller can be consistent with one or more of: Data Plane Development Kit (DPDK), OpenDataPlane (ODP), Infrastructure Programmer Development Kit (IPDK).

(23) FIG. 3 depicts an example system. Data in memory can be provided to a Network Interface. In some embodiments, one or more NICs can be coupled via a high speed interface to a switch that is on-circuit board with a GPU. Interface can include a silicon photonics or optical link, or other technology. Switch can be optionally utilized in some examples.

(24) FIG. 4 depicts an example system. In some examples, externally packaged NICs **408** can be coupled via a high speed interface **406** to switch **405** of connectivity **404**. In some examples, NICs **408** can be provided by a different manufacturer than that which manufactured GPU compute **102**. Switch **405** can be coupled to GPU compute **102** using a interface **402**. Interface **402** can include serial or parallel interfaces. Connectivity **404** can be communicatively coupled with NIC **408** using interface **406**. GPU compute **102** can be provided in a semiconductor package with interfaces **402**, connectivity **404** can be provided in a separate die than that of GPU compute **102**. Note that switch **405** can provide communication among GPU compute **102**, memory, **103**, and a memory device or

memory pool of external memory (e.g., dual inline memory modules (DIMMs) **412** via interface **410**. Interface **402**, interface **406**, and/or interface **410** can utilize one or more of the following standards and technology: Peripheral Component Interconnect express (PCIe), Compute Express Link (CXL), serializer de-serializer (SerDes), silicon photonics or optical interface, or other protocol or connectivity technology.

(25) Connectivity **404** can be provided in a same or different circuit board than that of GPU compute **102**. Connectivity **404** can be provided in a same or different package than that of GPU compute **102**. A package can include metal, plastic, glass, or ceramic casing that includes one or more semiconductor devices or integrated circuits. In some examples, NICs **408** can be provided on a same circuit board or different circuit board than that of GPU compute **102**. In some examples, NICs **408** can be provided on a same circuit board or different circuit board than that of connectivity **404**. In some examples, NICs **408** can be provided in a same or different semiconductor package than that of connectivity **404**. In some examples, NICs **408** can be provided in a same or different semiconductor package than that of GPU compute **102**. In some examples, NICs **408** can be provided in a same or different semiconductor package than that of GPU compute **102** and connectivity **404**. NICs **408** can provide communication among GPU compute **102**, memory **103**, memory **412**, and memory **414**, that may be accessible over the network or some interface. In some examples, memory **414** can be implemented as a memory pool of one or more DIMMs.

(26) FIG. 5 depicts an example of a GPU server. In some examples, scale out NICs **504-0** to **504-3** interface with GPU **502-0** to **502-7** via one or more PCIe switches **506-0** to **506-3**. Scale up NICs can be integrated into GPUs **502-0** to **502-7**, as described herein, to provide connectivity to on-circuit board GPUs. The scale out NICs can provide capability for a proprietary training cluster case and can be solely configured by CSPs, in accordance with the examples above. In some examples, the integrated NICs can be utilized for a GPU to communicate with other on-node devices (e.g., GPUs, memory, accelerators, and so forth) such as for scale-up connectivity, inter-node connectivity, or scale-out connectivity to communicate with other devices, whether the GPU and other devices are integrated into a same semiconductor device or different semiconductor devices. Integrated NICs can be controlled and configured by as CSP management entity via an interface to the CSP DC directly, or through another network interface device (e.g., a switch) or a network interface device (e.g., network interface device **210**).

(27) FIG. 6 depicts an example process. The process can be performed by a system with a GPU with one or more integrated or connected NICs for GPU-to-GPU connectivity. At **602**, a configuration of one or more NICs can be received. The configuration can be received from a CSP control plane software. The configuration can be received at a network interface device such as a smartNIC. In some examples, the configuration can specify one or more of: reliability protocol to apply, packet encryption to apply, packet decryption to apply, network protocol or frame format to utilize, and/or permitted destination for packets. At **604**, the one or more integrated or connected NICs can be configured by the configuration. For example, the data plane can be configured by the network interface device using the configuration. In some examples, the configuration is written in Protocol-independent Packet Processors (P4), although other formats can be utilized.

(28) FIGS. 7A-7D depict example GPU compute components. FIG. 7A illustrates a parallel processor **700**. The parallel processor **700** may be a GPU, GPGPU or the like as described herein. The various components of the parallel processor **700** may be implemented using one or more integrated circuit devices, such as programmable processors, application specific integrated circuits (ASICs), or field programmable gate arrays (FPGA).

(29) The parallel processor **700** includes a parallel processing unit **702**. The parallel processing unit includes an I/O unit **704** that enables communication with other devices, including other instances of the parallel processing unit **702**. The I/O unit **704** may be directly connected to other devices. For instance, the I/O unit **704** connects with other devices via the use of a hub or switch interface,

such as a memory hub. The connections between the memory hub **75** and the I/O unit **704** form a communication link. Within the parallel processing unit **702**, the I/O unit **704** connects with a host interface **706** and a memory crossbar **716**, where the host interface **706** receives commands directed to performing processing operations and the memory crossbar **716** receives commands directed to performing memory operations.

(30) When the host interface **706** receives a command buffer via the I/O unit **704**, the host interface **706** can direct work operations to perform those commands to a front end **708**. In one embodiment the front end **708** couples with a scheduler **710**, which is configured to distribute commands or other work items to a processing cluster array **712**. The scheduler **710** configures processing cluster array **712** is properly configured and in a valid state before tasks are distributed to the processing clusters of the processing cluster array **712**. The scheduler **710** may be implemented via firmware logic executing on a microcontroller. The microcontroller implemented scheduler **710** is configurable to perform complex scheduling and work distribution operations at coarse and fine granularity, enabling rapid preemption and context switching of threads executing on the processing cluster array **712**. Preferably, the host software can prove workloads for scheduling on the processing cluster array **712** via one of multiple graphics processing doorbells. In other examples, polling for new workloads or interrupts can be used to identify or indicate availability of work to perform. The workloads can then be automatically distributed across the processing cluster array **712** by the scheduler **710** logic within the scheduler microcontroller.

(31) The processing cluster array **712** can include up to “N” processing clusters (e.g., cluster **714A**, cluster **714B**, through cluster **714N**). At least one of cluster **714A-714N** of the processing cluster array **712** can execute a large number of concurrent threads. The scheduler **710** can allocate work to the clusters **714A-714N** of the processing cluster array **712** using various scheduling and/or work distribution algorithms, which may vary depending on the workload arising for a type of program or computation. The scheduling can be handled dynamically by the scheduler **710** or can be assisted in part by compiler logic during compilation of program logic configured for execution by the processing cluster array **712**. Optionally, different clusters **714A-714N** of the processing cluster array **712** can be allocated for processing different types of programs or for performing different types of computations.

(32) The processing cluster array **712** can be configured to perform various types of parallel processing operations. For example, the processing cluster array **712** is configured to perform general-purpose parallel compute operations. For example, the processing cluster array **712** can include logic to execute processing tasks including filtering of video and/or audio data, performing modeling operations, including physics operations, and performing data transformations.

(33) The processing cluster array **712** is configured to perform parallel graphics processing operations. In such embodiments in which the parallel processor **700** is configured to perform graphics processing operations, the processing cluster array **712** can include additional logic to support the execution of such graphics processing operations, including, but not limited to texture sampling logic to perform texture operations, as well as tessellation logic and other vertex processing logic. Additionally, the processing cluster array **712** can be configured to execute graphics processing related shader programs such as, but not limited to vertex shaders, tessellation shaders, geometry shaders, and pixel shaders. The parallel processing unit **702** can transfer data from system memory via the I/O unit **704** for processing. During processing the transferred data can be stored to on-chip memory (e.g., parallel processor memory **722**) during processing, then written back to system memory.

(34) In embodiments in which the parallel processing unit **702** is used to perform graphics processing, the scheduler **710** may be configured to divide the processing workload into approximately equal sized tasks, to better enable distribution of the graphics processing operations to multiple clusters **714A-714N** of the processing cluster array **712**. In some of these embodiments, portions of the processing cluster array **712** can be configured to perform different types of

processing. For example, a first portion may be configured to perform vertex shading and topology generation, a second portion may be configured to perform tessellation and geometry shading, and a third portion may be configured to perform pixel shading or other screen space operations, to produce a rendered image for display. Intermediate data produced by one or more of the clusters **714A-714N** may be stored in buffers to allow the intermediate data to be transmitted between clusters **714A-714N** for further processing.

(35) During operation, the processing cluster array **712** can receive processing tasks to be executed via the scheduler **710**, which receives commands defining processing tasks from front end **708**. For graphics processing operations, processing tasks can include indices of data to be processed, e.g., surface (patch) data, primitive data, vertex data, and/or pixel data, as well as state parameters and commands defining how the data is to be processed (e.g., what program is to be executed). The scheduler **710** may be configured to fetch the indices corresponding to the tasks or may receive the indices from the front end **708**. The front end **708** can configure the processing cluster array **712** to a valid state before the workload specified by incoming command buffers (e.g., batch-buffers, push buffers, etc.) is initiated.

(36) At least one of the one or more instances of the parallel processing unit **702** can couple with parallel processor memory **722**. The parallel processor memory **722** can be accessed via the memory crossbar **716**, which can receive memory requests from the processing cluster array **712** as well as the I/O unit **704**. The memory crossbar **716** can access the parallel processor memory **722** via a memory interface **718**. The memory interface **718** can include multiple partition units (e.g., partition unit **720A**, partition unit **720B**, through partition unit **720N**) that can couple to a portion (e.g., memory unit) of parallel processor memory **722**. The number of partition units **720A-720N** may be configured to be equal to the number of memory units, such that a first partition unit **720A** has a corresponding first memory unit **724A**, a second partition unit **720B** has a corresponding second memory unit **724B**, and an Nth partition unit **720N** has a corresponding Nth memory unit **724N**. In other embodiments, the number of partition units **720A-720N** may not be equal to the number of memory devices.

(37) The memory units **724A-724N** can include various types of memory devices, including dynamic random-access memory (DRAM) or graphics random access memory, such as synchronous graphics random access memory (SGRAM), including graphics double data rate (GDDR) memory. Optionally, the memory units **724A-724N** may also include 3D stacked memory, including but not limited to high bandwidth memory (HBM). Persons skilled in the art will appreciate that the specific implementation of the memory units **724A-724N** can vary and can be selected from one of various conventional designs. Render targets, such as frame buffers or texture maps may be stored across the memory units **724A-724N**, allowing partition units **720A-720N** to write portions of a render target in parallel to efficiently use the available bandwidth of parallel processor memory **722**. In some embodiments, a local instance of the parallel processor memory **722** may be excluded in favor of a unified memory design that utilizes system memory in conjunction with local cache memory.

(38) Optionally, any one of the clusters **714A-714N** of the processing cluster array **712** has the ability to process data that will be written to any of the memory units **724A-724N** within parallel processor memory **722**. The memory crossbar **716** can be configured to transfer the output of at least one of cluster **714A-714N** to any partition unit **720A-720N** or to another cluster **714A-714N**, which can perform additional processing operations on the output. At least one of cluster **714A-714N** can communicate with the memory interface **718** through the memory crossbar **716** to read from or write to various external memory devices. In one of the embodiments with the memory crossbar **716** the memory crossbar **716** has a connection to the memory interface **718** to communicate with the I/O unit **704**, as well as a connection to a local instance of the parallel processor memory **722**, enabling the processing units within the different processing clusters **714A-714N** to communicate with system memory or other memory that is not local to the parallel

processing unit **702**. Generally, the memory crossbar **716** may, for example, be able to use virtual channels to separate traffic streams between the clusters **714A-714N** and the partition units **720A-720N**.

(39) While a single instance of the parallel processing unit **702** is illustrated within the parallel processor **700**, any number of instances of the parallel processing unit **702** can be included. For example, multiple instances of the parallel processing unit **702** can be provided on a single add-in card, or multiple add-in cards can be interconnected. For example, the parallel processor **700** can be an add-in device, which may be a graphics card such as a discrete graphics card that includes one or more GPUs, one or more memory devices, and device-to-device or network or fabric interfaces. The different instances of the parallel processing unit **702** can be configured to inter-operate even if the different instances have different numbers of processing cores, different amounts of local parallel processor memory, and/or other configuration differences. Optionally, some instances of the parallel processing unit **702** can include higher precision floating point units relative to other instances. Systems incorporating one or more instances of the parallel processing unit **702** or the parallel processor **700** can be implemented in a variety of configurations and form factors, including but not limited to desktop, laptop, or handheld personal computers, servers, workstations, game consoles, and/or embedded systems. An orchestrator can form composite nodes for workload performance using one or more of: disaggregated processor resources, cache resources, memory resources, storage resources, and networking resources.

(40) FIG. 7B is a block diagram of a partition unit **720**. The partition unit **720** may be an instance of one of the partition units **720A-720N** of FIG. 7A. As illustrated, the partition unit **720** includes an L2 cache **721**, a frame buffer interface **725**, and a ROP **726** (raster operations unit). The L2 cache **721** is a read/write cache that is configured to perform load and store operations received from the memory crossbar **716** and ROP **726**. Read misses and urgent write-back requests are output by L2 cache **721** to frame buffer interface **725** for processing. Updates can also be sent to the frame buffer via the frame buffer interface **725** for processing. In one embodiment the frame buffer interface **725** interfaces with one of the memory units in parallel processor memory, such as the memory units **724A-724N** of FIG. 7A (e.g., within parallel processor memory **722**). The partition unit **720** may additionally or alternatively also interface with one of the memory units in parallel processor memory via a memory controller (not shown).

(41) In graphics applications, the ROP **726** is a processing unit that performs raster operations such as stencil, z test, blending, and the like. The ROP **726** then outputs processed graphics data that is stored in graphics memory. In some embodiments the ROP **726** includes or couples with a CODEC **727** that includes compression logic to compress depth or color data that is written to memory or the L2 cache **721** and decompress depth or color data that is read from memory or the L2 cache **721**. The compression logic can be lossless compression logic that makes use of one or more of multiple compression algorithms. The type of compression that is performed by the CODEC **727** can vary based on the statistical characteristics of the data to be compressed. For example, in one embodiment, delta color compression is performed on depth and color data on a per-tile basis. In one embodiment the CODEC **727** includes compression and decompression logic that can compress and decompress compute data associated with machine learning operations. The CODEC **727** can, for example, compress sparse matrix data for sparse machine learning operations. The CODEC **727** can also compress sparse matrix data that is encoded in a sparse matrix format (e.g., coordinate list encoding (COO), compressed sparse row (CSR), compressed sparse column (CSC), etc.) to generate compressed and encoded sparse matrix data. The compressed and encoded sparse matrix data can be decompressed and/or decoded before being processed by processing elements or the processing elements can be configured to consume compressed, encoded, or compressed and encoded data for processing.

(42) The ROP **726** may be included within at least one processing cluster (e.g., cluster **714A-714N** of FIG. 7A) instead of within the partition unit **720**. In such embodiment, read and write requests

for pixel data are transmitted over the memory crossbar **716** instead of pixel fragment data. The processed graphics data may be displayed on a display device, such as one of the one or more display device(s), routed for further processing by processor(s), or routed for further processing by one of the processing entities within a parallel processor **700**.

(43) FIG. 7C is a block diagram of a processing cluster **714** within a parallel processing unit. For example, the processing cluster is an instance of one of the processing clusters **714A-714N** of FIG. 7A. The processing cluster **714** can be configured to execute many threads in parallel, where the term “thread” refers to an instance of a particular program executing on a particular set of input data. Optionally, single-instruction, multiple-data (SIMD) instruction issue techniques may be used to support parallel execution of a large number of threads without providing multiple independent instruction units. Alternatively, single-instruction, multiple-thread (SIMT) techniques may be used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of processing engines within at least one of the processing clusters. Unlike a SIMD execution regime, where all processing engines typically execute identical instructions, SIMT execution allows different threads to more readily follow divergent execution paths through a given thread program. Persons skilled in the art will understand that a SIMD processing regime represents a functional subset of a SIMT processing regime.

(44) Operation of the processing cluster **714** can be controlled via a pipeline manager **732** that distributes processing tasks to SIMT parallel processors. The pipeline manager **732** receives instructions from the scheduler **710** of FIG. 7A and manages execution of those instructions via a graphics multiprocessor **734** and/or a texture unit **736**. The illustrated graphics multiprocessor **734** is an exemplary instance of a SIMT parallel processor. However, various types of SIMT parallel processors of differing architectures may be included within the processing cluster **714**. One or more instances of the graphics multiprocessor **734** can be included within a processing cluster **714**. The graphics multiprocessor **734** can process data and a data crossbar **740** can be used to distribute the processed data to one of multiple possible destinations, including other shader units. The pipeline manager **732** can facilitate the distribution of processed data by specifying destinations for processed data to be distributed via the data crossbar **740**.

(45) At least one of graphics multiprocessor **734** within the processing cluster **714** can include an identical set of functional execution logic (e.g., arithmetic logic units, load-store units, etc.). The functional execution logic can be configured in a pipelined manner in which new instructions can be issued before previous instructions are complete. The functional execution logic supports a variety of operations including integer and floating-point arithmetic, comparison operations, Boolean operations, bit-shifting, and computation of various algebraic functions. The same functional-unit hardware could be leveraged to perform different operations and any combination of functional units may be present.

(46) The instructions transmitted to the processing cluster **714** constitute a thread. A set of threads executing across the set of parallel processing engines is a thread group. A thread group executes the same program on different input data. At least one thread within a thread group can be assigned to a different processing engine within a graphics multiprocessor **734**. A thread group may include fewer threads than the number of processing engines within the graphics multiprocessor **734**. When a thread group includes fewer threads than the number of processing engines, one or more of the processing engines may be idle during cycles in which that thread group is being processed. A thread group may also include more threads than the number of processing engines within the graphics multiprocessor **734**. When the thread group includes more threads than the number of processing engines within the graphics multiprocessor **734**, processing can be performed over consecutive clock cycles. Optionally, multiple thread groups can be executed concurrently on the graphics multiprocessor **734**.

(47) The graphics multiprocessor **734** may include an internal cache memory to perform load and

store operations. Optionally, the graphics multiprocessor **734** can forego an internal cache and use a cache memory (e.g., level 1 (L1) cache **748**) within the processing cluster **714**. At least one graphics multiprocessor **734** also has access to level 2 (L2) caches within the partition units (e.g., partition units **720A-720N** of FIG. 7A) that are shared among all processing clusters **714** and may be used to transfer data between threads. The graphics multiprocessor **734** may also access off-chip global memory, which can include one or more of local parallel processor memory and/or system memory. Any memory external to the parallel processing unit **702** may be used as global memory. Embodiments in which the processing cluster **714** includes multiple instances of the graphics multiprocessor **734** can share common instructions and data, which may be stored in the L1 cache **748**.

(48) At least one processing cluster **714** may include an MMU **745** (memory management unit) that is configured to map virtual addresses into physical addresses. In other embodiments, one or more instances of the MMU **745** may reside within the memory interface **718** of FIG. 7A. The MMU **745** includes a set of page table entries (PTEs) used to map a virtual address to a physical address of a tile and optionally a cache line index. The MMU **745** may include address translation lookaside buffers (TLB) or caches that may reside within the graphics multiprocessor **734** or the L1 cache **748** of processing cluster **714**. The physical address is processed to distribute surface data access locality to allow efficient request interleaving among partition units. The cache line index may be used to determine whether a request for a cache line is a hit or miss.

(49) In graphics and computing applications, a processing cluster **714** may be configured such that at least one graphics multiprocessor **734** is coupled to a texture unit **736** for performing texture mapping operations, e.g., determining texture sample positions, reading texture data, and filtering the texture data. Texture data is read from an internal texture L1 cache (not shown) or in some embodiments from the L1 cache within graphics multiprocessor **734** and is fetched from an L2 cache, local parallel processor memory, or system memory, as needed. At least one graphics multiprocessor **734** outputs processed tasks to the data crossbar **740** to provide the processed task to another processing cluster **714** for further processing or to store the processed task in an L2 cache, local parallel processor memory, or system memory via the memory crossbar **716**. A preROP **742** (pre-raster operations unit) is configured to receive data from graphics multiprocessor **734**, direct data to ROP units, which may be located with partition units as described herein (e.g., partition units **720A-720N** of FIG. 7A). The preROP **742** unit can perform optimizations for color blending, organize pixel color data, and perform address translations.

(50) It will be appreciated that the core architecture described herein is illustrative and that variations and modifications are possible. Any number of processing units, e.g., graphics multiprocessor **734**, texture units **736**, preROPs **742**, etc., may be included within a processing cluster **714**. Further, while only one processing cluster **714** is shown, a parallel processing unit as described herein may include any number of instances of the processing cluster **714**. Optionally, at least one processing cluster **714** can be configured to operate independently of other processing clusters **714** using separate and distinct processing units, L1 caches, L2 caches, etc.

(51) FIG. 7D shows an example of the graphics multiprocessor **734** in which the graphics multiprocessor **734** couples with the pipeline manager **732** of the processing cluster **714**. The graphics multiprocessor **734** has an execution pipeline including but not limited to an instruction cache **752**, an instruction unit **754**, an address mapping unit **756**, a register file **758**, one or more general purpose graphics processing unit (GPGPU) cores **762**, and one or more load/store units **766**. The GPGPU cores **762** and load/store units **766** are coupled with cache memory **772** and shared memory **770** via a memory and cache interconnect **768**. The graphics multiprocessor **734** may additionally include tensor and/or ray-tracing cores **763** that include hardware logic to accelerate matrix and/or ray-tracing operations.

(52) The instruction cache **752** may receive a stream of instructions to execute from the pipeline manager **732**. The instructions are cached in the instruction cache **752** and dispatched for execution

by the instruction unit **754**. The instruction unit **754** can dispatch instructions as thread groups (e.g., warps), with at least one thread of the thread group assigned to a different execution unit within GPGPU core **762**. An instruction can access a local, shared, or global address space by specifying an address within a unified address space. The address mapping unit **756** can be used to translate addresses in the unified address space into a distinct memory address that can be accessed by the load/store units **766**.

(53) The register file **758** provides a set of registers for the functional units of the graphics multiprocessor **734**. The register file **758** provides temporary storage for operands connected to the data paths of the functional units (e.g., GPGPU cores **762**, load/store units **766**) of the graphics multiprocessor **734**. The register file **758** may be divided between at least one of the functional units such that at least one functional unit is allocated a dedicated portion of the register file **758**. For example, the register file **758** may be divided between the different warps being executed by the graphics multiprocessor **734**.

(54) The GPGPU cores **762** can include floating point units (FPUs) and/or integer arithmetic logic units (ALUs) that are used to execute instructions of the graphics multiprocessor **734**. In some implementations, the GPGPU cores **762** can include hardware logic that may otherwise reside within the tensor and/or ray-tracing cores **763**. The GPGPU cores **762** can be similar in architecture or can differ in architecture. For example and in one embodiment, a first portion of the GPGPU cores **762** include a single precision FPU and an integer ALU while a second portion of the GPGPU cores include a double precision FPU. Optionally, the FPUs can implement the IEEE 754-2008 standard for floating point arithmetic or enable variable precision floating point arithmetic. The graphics multiprocessor **734** can additionally include one or more fixed function or special function units to perform specific functions such as copy rectangle or pixel blending operations. One or more of the GPGPU cores can also include fixed or special function logic.

(55) The GPGPU cores **762** may include SIMD logic capable of performing a single instruction on multiple sets of data. Optionally, GPGPU cores **762** can physically execute SIMD4, SIMD8, and SIMD16 instructions and logically execute SIMD1, SIMD2, and SIMD3 instructions. The SIMD instructions for the GPGPU cores can be generated at compile time by a shader compiler or automatically generated when executing programs written and compiled for single program multiple data (SPMD) or SIMT architectures. Multiple threads of a program configured for the SIMT execution model can be executed via a single SIMD instruction. For example and in one embodiment, eight SIMT threads that perform the same or similar operations can be executed in parallel via a single SIMD8 logic unit.

(56) The memory and cache interconnect **768** is an interconnect network that connects at least one of the functional units of the graphics multiprocessor **734** to the register file **758** and to the shared memory **770**. For example, the memory and cache interconnect **768** is a crossbar interconnect that allows the load/store unit **766** to implement load and store operations between the shared memory **770** and the register file **758**. The register file **758** can operate at the same frequency as the GPGPU cores **762**, thus data transfer between the GPGPU cores **762** and the register file **758** is very low latency. The shared memory **770** can be used to enable communication between threads that execute on the functional units within the graphics multiprocessor **734**. The cache memory **772** can be used as a data cache for example, to cache texture data communicated between the functional units and the texture unit **736**. The shared memory **770** can also be used as a program managed cached. The shared memory **770** and the cache memory **772** can couple with the data crossbar **740** to enable communication with other components of the processing cluster. Threads executing on the GPGPU cores **762** can programmatically store data within the shared memory in addition to the automatically cached data that is stored within the cache memory **772**.

(57) FIG. **8** depicts an example network interface device. Various hardware and software resources in the network interface can be configured by another network interface device with CSP configurations, as described herein. In some examples, network interface **800** can be implemented

as a network interface controller, network interface card, a host fabric interface (HFI), or host bus adapter (HBA), and such examples can be interchangeable. Network interface **800** can be coupled to one or more servers using a bus, PCIe, CXL, or DDR. Network interface **800** may be embodied as part of a system-on-a-chip (SoC) that includes one or more processors, or included on a multichip package that also contains one or more processors.

(58) Some examples of network device **800** are part of an Infrastructure Processing Unit (IPU) or data processing unit (DPU) or utilized by an IPU or DPU. An xPU can refer at least to an IPU, DPU, GPU, GPGPU, or other processing units (e.g., accelerator devices). An IPU or DPU can include a network interface with one or more programmable pipelines or fixed function processors to perform offload of operations that could have been performed by a CPU. The IPU or DPU can include one or more memory devices. In some examples, the IPU or DPU can perform virtual switch operations, manage storage transactions (e.g., compression, cryptography, virtualization), and manage operations performed on other IPUs, DPUs, servers, or devices.

(59) Network interface **800** can include transceiver **802**, processors **804**, transmit queue **806**, receive queue **808**, memory **810**, and bus interface **812**, and DMA engine **852**. Transceiver **802** can be capable of receiving and transmitting packets in conformance with the applicable protocols such as Ethernet as described in IEEE 802.3, although other protocols may be used. Transceiver **802** can receive and transmit packets from and to a network via a network medium (not depicted).

Transceiver **802** can include PHY circuitry **814** and media access control (MAC) circuitry **816**. PHY circuitry **814** can include encoding and decoding circuitry (not shown) to encode and decode data packets according to applicable physical layer specifications or standards. MAC circuitry **816** can be configured to perform MAC address filtering on received packets, process MAC headers of received packets by verifying data integrity, remove preambles and padding, and provide packet content for processing by higher layers. MAC circuitry **816** can be configured to assemble data to be transmitted into packets, that include destination and source addresses along with network control information and error detection hash values.

(60) Processors **804** can be any a combination of: a processor, core, graphics processing unit (GPU), field programmable gate array (FPGA), application specific integrated circuit (ASIC), or other programmable hardware device that allow programming of network interface **800**. For example, a “smart network interface” or SmartNIC can provide packet processing capabilities in the network interface using processors **804**.

(61) Processors **804** can include a programmable processing pipeline that is programmable by one or more of: Protocol-independent Packet Processors (P4), Software for Open Networking in the Cloud (SONiC), Broadcom® Network Programming Language (NPL), NVIDIA® CUDA®, NVIDIA® DOCA™, Infrastructure Programmer Development Kit (IPDK), or x86 compatible executable binaries or other executable binaries. A programmable processing pipeline can include one or more match-action units (MAUs) that can schedule packets for transmission using one or multiple granularity lists, as described herein. Processors, FPGAs, other specialized processors, controllers, devices, and/or circuits can be used utilized for packet processing or packet modification. Ternary content-addressable memory (TCAM) can be used for parallel match-action or look-up operations on packet header content.

(62) Packet allocator **824** can provide distribution of received packets for processing by multiple CPUs or cores using receive side scaling (RSS). When packet allocator **824** uses RSS, packet allocator **824** can calculate a hash or make another determination based on contents of a received packet to determine which CPU or core is to process a packet.

(63) Interrupt coalesce **822** can perform interrupt moderation whereby network interface interrupt coalesce **822** waits for multiple packets to arrive, or for a time-out to expire, before generating an interrupt to host system to process received packet(s). Receive Segment Coalescing (RSC) can be performed by network interface **800** whereby portions of incoming packets are combined into segments of a packet. Network interface **800** provides this coalesced packet to an application.

(64) Direct memory access (DMA) engine **852** can copy a packet header, packet payload, and/or descriptor directly from host memory to the network interface or vice versa, instead of copying the packet to an intermediate buffer at the host and then using another copy operation from the intermediate buffer to the destination buffer.

(65) Memory **810** can be any type of volatile or non-volatile memory device and can store any queue or instructions used to program network interface **800**. Transmit queue **806** can include data or references to data for transmission by network interface. Receive queue **808** can include data or references to data that was received by network interface from a network. Descriptor queues **820** can include descriptors that reference data or packets in transmit queue **806** or receive queue **808**. Bus interface **812** can provide an interface with host device (not depicted). For example, bus interface **812** can be compatible with or based at least in part on PCI, PCI Express, PCI-x, Serial ATA, and/or USB (although other interconnection standards may be used), or proprietary variations thereof.

(66) FIG. 9 depicts an example computing system. Components of system **900** can configure NICs of a GPU with integrated or connected NICs. Components of system **900** can utilize a GPU with integrated or connected NICs for GPU-to-GPU communication, as described herein. System **900** includes processor **910**, which provides processing, operation management, and execution of instructions for system **900**. Processor **910** can include any type of microprocessor, central processing unit (CPU), graphics processing unit (GPU), processing core, or other processing hardware to provide processing for system **900**, or a combination of processors. In some examples, a CPU and GPU can be integrated into a single package. Processor **910** controls the overall operation of system **900**, and can be or include, one or more programmable general-purpose or special-purpose microprocessors, digital signal processors (DSPs), programmable controllers, application specific integrated circuits (ASICs), programmable logic devices (PLDs), or the like, or a combination of such devices.

(67) In one example, system **900** includes interface **912** coupled to processor **910**, which can represent a higher speed interface or a high throughput interface for system components that needs higher bandwidth connections, such as memory subsystem **920** or graphics interface components **940**, or accelerators **942**. Interface **912** represents an interface circuit, which can be a standalone component or integrated onto a processor die. Where present, graphics interface **940** interfaces to graphics components for providing a visual display to a user of system **900**. In one example, graphics interface **940** can drive a high definition (HD) display that provides an output to a user. High definition can refer to a display having a pixel density of approximately 100 PPI (pixels per inch) or greater and can include formats such as full HD (e.g., 1080p), retina displays, 4K (ultra-high definition or UHD), or others. In one example, the display can include a touchscreen display. In one example, graphics interface **940** generates a display based on data stored in memory **930** or based on operations executed by processor **910** or both. In one example, graphics interface **940** generates a display based on data stored in memory **930** or based on operations executed by processor **910** or both.

(68) Accelerators **942** can be a fixed function or programmable offload engine that can be accessed or used by a processor **910**. For example, an accelerator among accelerators **942** can provide compression (DC) capability, cryptography services such as public key encryption (PKE), cipher, hash/authentication capabilities, decryption, or other capabilities or services. In some embodiments, in addition or alternatively, an accelerator among accelerators **942** provides field select controller capabilities as described herein. In some cases, accelerators **942** can be integrated into a CPU socket (e.g., a connector to a motherboard or circuit board that includes a CPU and provides an electrical interface with the CPU). For example, accelerators **942** can include a single or multi-core processor, graphics processing unit, logical execution unit single or multi-level cache, functional units usable to independently execute programs or threads, application specific integrated circuits (ASICs), neural network processors (NNPs), programmable control logic, and programmable

processing elements such as field programmable gate arrays (FPGAs) or programmable logic devices (PLDs). Accelerators **942** can provide multiple neural networks, CPUs, processor cores, general purpose graphics processing units, or graphics processing units can be made available for use by artificial intelligence (AI) or machine learning (ML) models. For example, the AI model can use or include one or more of: a reinforcement learning scheme, Q-learning scheme, deep-Q learning, or Asynchronous Advantage Actor-Critic (A3C), combinatorial neural network, recurrent combinatorial neural network, or other AI or ML model. Multiple neural networks, processor cores, or graphics processing units can be made available for use by AI or ML models.

(69) Memory subsystem **920** represents the main memory of system **900** and provides storage for code to be executed by processor **910**, or data values to be used in executing a routine. Memory subsystem **920** can include one or more memory devices **930** such as read-only memory (ROM), flash memory, one or more varieties of random access memory (RAM) such as DRAM, or other memory devices, or a combination of such devices. Memory **930** stores and hosts, among other things, operating system (OS) **932** to provide a software platform for execution of instructions in system **900**. Additionally, applications **934** can execute on the software platform of OS **932** from memory **930**. Applications **934** represent programs that have their own operational logic to perform execution of one or more functions. Processes **936** represent agents or routines that provide auxiliary functions to OS **932** or one or more applications **934** or a combination. OS **932**, applications **934**, and processes **936** provide software logic to provide functions for system **900**. In some examples, applications **934** can request GPU-to-GPU communication using integrated or connected NICs, as described herein.

(70) In one example, memory subsystem **920** includes memory controller **922**, which is a memory controller to generate and issue commands to memory **930**. It will be understood that memory controller **922** could be a physical part of processor **910** or a physical part of interface **912**. For example, memory controller **922** can be an integrated memory controller, integrated onto a circuit with processor **910**.

(71) In some examples, OS **932** can be Linux®, Windows® Server or personal computer, FreeBSD®, Android®, MacOS®, iOS®, VMware vSphere, openSUSE, RHEL, CentOS, Debian, Ubuntu, or any other operating system. The OS and driver can execute on a CPU sold or designed by Intel®, ARM®, AMD®, NVIDIA®, Broadcom®, Qualcomm®, IBM®, Texas Instruments®, among others.

(72) While not specifically illustrated, it will be understood that system **900** can include one or more buses or bus systems between devices, such as a memory bus, a graphics bus, interface buses, or others. Buses or other signal lines can communicatively or electrically couple components together, or both communicatively and electrically couple the components. Buses can include physical communication lines, point-to-point connections, bridges, adapters, controllers, or other circuitry or a combination. Buses can include, for example, one or more of a system bus, a Peripheral Component Interconnect (PCI) bus, a Hyper Transport or industry standard architecture (ISA) bus, a small computer system interface (SCSI) bus, a universal serial bus (USB), or an Institute of Electrical and Electronics Engineers (IEEE) standard 1394 bus (Firewire).

(73) In one example, system **900** includes interface **914**, which can be coupled to interface **912**. In one example, interface **914** represents an interface circuit, which can include standalone components and integrated circuitry. In one example, multiple user interface components or peripheral components, or both, couple to interface **914**. Network interface **950** provides system **900** the ability to communicate with remote devices (e.g., servers or other computing devices) over one or more networks. Network interface **950** can include an Ethernet adapter, wireless interconnection components, cellular network interconnection components, USB (universal serial bus), or other wired or wireless standards-based or proprietary interfaces. Network interface **950** can transmit data to a device that is in the same data center or rack or a remote device, which can include sending data stored in memory. Network interface **950** (e.g., packet processing device) can

execute a virtual switch to provide virtual machine-to-virtual machine communications for virtual machines (or containers, microservices, or other virtual execution environments) in a same server or among different servers.

(74) Some examples of network interface **950** are part of an Infrastructure Processing Unit (IPU) or data processing unit (DPU) or utilized by an IPU or DPU. An xPU can refer at least to an IPU, DPU, GPU, GPGPU, or other processing units (e.g., accelerator devices). An IPU or DPU can include a network interface with one or more programmable pipelines or fixed function processors to perform offload of operations that could have been performed by a CPU. The IPU or DPU can include one or more memory devices. In some examples, the IPU or DPU can perform virtual switch operations, manage storage transactions (e.g., compression, cryptography, virtualization), and manage operations performed on other IPUs, DPUs, servers, or devices.

(75) Graphics **940**, accelerators **942**, and/or network interface **950** can include Direct memory access (DMA) circuitry that can copy data to or from host memory, instead of copying the packet to an intermediate buffer at the host and then using another copy operation from the intermediate buffer to the destination buffer.

(76) In one example, system **900** includes one or more input/output (I/O) interface(s) **960**. I/O interface **960** can include one or more interface components through which a user interacts with system **900** (e.g., audio, alphanumeric, tactile/touch, or other interfacing). Peripheral interface **970** can include any hardware interface not specifically mentioned above. Peripherals refer generally to devices that connect dependently to system **900**. A dependent connection is one where system **900** provides the software platform or hardware platform or both on which operation executes, and with which a user interacts.

(77) In one example, system **900** includes storage subsystem **980** to store data in a nonvolatile manner. In one example, in certain system implementations, at least certain components of storage **980** can overlap with components of memory subsystem **920**. Storage subsystem **980** includes storage device(s) **984**, which can be or include any conventional medium for storing large amounts of data in a nonvolatile manner, such as one or more magnetic, solid state, or optical based disks, or a combination. Storage **984** holds code or instructions and data **986** in a persistent state (e.g., the value is retained despite interruption of power to system **900**). Storage **984** can be generically considered to be a “memory,” although memory **930** is typically the executing or operating memory to provide instructions to processor **910**. Whereas storage **984** is nonvolatile, memory **930** can include volatile memory (e.g., the value or state of the data is indeterminate if power is interrupted to system **900**). In one example, storage subsystem **980** includes controller **982** to interface with storage **984**. In one example controller **982** is a physical part of interface **914** or processor **910** or can include circuits or logic in both processor **910** and interface **914**.

(78) A volatile memory is memory whose state (and therefore the data stored in it) is indeterminate if power is interrupted to the device. Dynamic volatile memory requires refreshing the data stored in the device to maintain state. One example of dynamic volatile memory includes DRAM (Dynamic Random Access Memory), or some variant such as Synchronous DRAM (SDRAM). Another example of volatile memory includes cache or static random access memory (SRAM).

(79) A non-volatile memory (NVM) device is a memory whose state is determinate even if power is interrupted to the device. In one embodiment, the NVM device can comprise a block addressable memory device, such as NAND technologies, or more specifically, multi-threshold level NAND flash memory (for example, Single-Level Cell (“SLC”), Multi-Level Cell (“MLC”), Quad-Level Cell (“QLC”), Tri-Level Cell (“TLC”), or some other NAND). A NVM device can also comprise a byte-addressable write-in-place three dimensional cross point memory device, or other byte addressable write-in-place NVM device (also referred to as persistent memory), such as single or multi-level Phase Change Memory (PCM) or phase change memory with a switch (PCMS), Intel® Optane™ memory, or NVM devices that use chalcogenide phase change material (for example, chalcogenide glass).

(80) A power source (not depicted) provides power to the components of system **900**. More specifically, power source typically interfaces to one or multiple power supplies in system **900** to provide power to the components of system **900**. In one example, the power supply includes an AC to DC (alternating current to direct current) adapter to plug into a wall outlet. Such AC power can be renewable energy (e.g., solar power) power source. In one example, power source includes a DC power source, such as an external AC to DC converter. In one example, power source or power supply includes wireless charging hardware to charge via proximity to a charging field. In one example, power source can include an internal battery, alternating current supply, motion-based power supply, solar power supply, or fuel cell source.

(81) In an example, system **900** can be implemented using interconnected compute sleds of processors, memories, storages, network interfaces, and other components. High speed interconnects can be used such as: Ethernet (IEEE 802.3), remote direct memory access (RDMA), InfiniBand, Internet Wide Area RDMA Protocol (iWARP), Transmission Control Protocol (TCP), User Datagram Protocol (UDP), quick UDP Internet Connections (QUIC), RDMA over Converged Ethernet (RoCE), Peripheral Component Interconnect express (PCIe), Intel QuickPath Interconnect (QPI), Intel Ultra Path Interconnect (UPI), Intel On-Chip System Fabric (IOSF), Omni-Path, Compute Express Link (CXL), HyperTransport, high-speed fabric, NVLink, Advanced Microcontroller Bus Architecture (AMB A) interconnect, OpenCAPI, Gen-Z, Infinity Fabric (IF), Cache Coherent Interconnect for Accelerators (COX), 3GPP Long Term Evolution (LTE) (4G), 3GPP 5G, and variations thereof. Data can be copied or stored to virtualized storage nodes or accessed using a protocol such as NVMe over Fabrics (NVMe-oF) or NVMe.

(82) Communications between devices can take place using a network that provides die-to-die communications; chip-to-chip communications; circuit board-to-circuit board communications; and/or package-to-package communications.

(83) FIG. **10** depicts an example system. In this system, IPU **1000** manages performance of one or more processes using one or more of processors **1006**, processors **1010**, accelerators **1020**, memory pool **1030**, or servers **1040-0** to **1040-N**, where N is an integer of 1 or more. In some examples, network interface device **210** can be implemented to include at least some components of IPU **1000**. In some examples, processors **1006** of IPU **1000** can execute one or more processes, applications, VMs, containers, microservices, and so forth that request performance of workloads by one or more of: processors **1010**, accelerators **1020**, memory pool **1030**, and/or servers **1040-0** to **1040-N**. IPU **1000** can utilize network interface **1002** or one or more device interfaces to communicate with processors **1010**, accelerators **1020**, memory pool **1030**, and/or servers **1040-0** to **1040-N**. IPU **1000** can utilize programmable pipeline **1004** to process packets that are to be transmitted from network interface **1002** or packets received from network interface **1002**.

Programmable pipeline **1004** and/or processors **1006** can be configured to perform configuration of data planes of NICs connected to or integrated into GPUs, as described herein.

(84) Embodiments herein may be implemented in various types of computing, smart phones, tablets, personal computers, and networking equipment, such as switches, routers, racks, and blade servers such as those employed in a data center and/or server farm environment. The servers used in data centers and server farms comprise arrayed server configurations such as rack-based servers or blade servers. These servers are interconnected in communication via various network provisions, such as partitioning sets of servers into Local Area Networks (LANs) with appropriate switching and routing facilities between the LANs to form a private Intranet. For example, cloud hosting facilities may typically employ large data centers with a multitude of servers. A blade comprises a separate computing platform that is configured to perform server-type functions, that is, a “server on a card.” Accordingly, a blade includes components common to conventional servers, including a main printed circuit board (main board) providing internal wiring (e.g., buses) for coupling appropriate integrated circuits (ICs) and other components mounted to the board.

(85) In some examples, network interface and other embodiments described herein can be used in

connection with a base station (e.g., 3G, 4G, 5G and so forth), macro base station (e.g., 5G networks), picostation (e.g., an IEEE 802.11 compatible access point), nanostation (e.g., for Point-to-MultiPoint (PtMP) applications), on-premises data centers, off-premises data centers, edge network elements, fog network elements, and/or hybrid data centers (e.g., data center that use virtualization, cloud and software-defined networking to deliver application workloads across physical data centers and distributed multi-cloud environments).

(86) Various examples may be implemented using hardware elements, software elements, or a combination of both. In some examples, hardware elements may include devices, components, processors, microprocessors, circuits, circuit elements (e.g., transistors, resistors, capacitors, inductors, and so forth), integrated circuits, ASICs, PLDs, DSPs, FPGAs, memory units, logic gates, registers, semiconductor device, chips, microchips, chip sets, and so forth. In some examples, software elements may include software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, APIs, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or any combination thereof. Determining whether an example is implemented using hardware elements and/or software elements may vary in accordance with any number of factors, such as desired computational rate, power levels, heat tolerances, processing cycle budget, input data rates, output data rates, memory resources, data bus speeds and other design or performance constraints, as desired for a given implementation. A processor can be one or more combination of a hardware state machine, digital control logic, central processing unit, or any hardware, firmware and/or software elements.

(87) Some examples may be implemented using or as an article of manufacture or at least one computer-readable medium. A computer-readable medium may include a non-transitory storage medium to store logic. In some examples, the non-transitory storage medium may include one or more types of computer-readable storage media capable of storing electronic data, including volatile memory or non-volatile memory, removable or non-removable memory, erasable or non-erasable memory, writable or re-writable memory, and so forth. In some examples, the logic may include various software elements, such as software components, programs, applications, computer programs, application programs, system programs, machine programs, operating system software, middleware, firmware, software modules, routines, subroutines, functions, methods, procedures, software interfaces, API, instruction sets, computing code, computer code, code segments, computer code segments, words, values, symbols, or any combination thereof.

(88) According to some examples, a computer-readable medium may include a non-transitory storage medium to store or maintain instructions that when executed by a machine, computing device or system, cause the machine, computing device or system to perform methods and/or operations in accordance with the described examples. The instructions may include any suitable type of code, such as source code, compiled code, interpreted code, executable code, static code, dynamic code, and the like. The instructions may be implemented according to a predefined computer language, manner or syntax, for instructing a machine, computing device or system to perform a certain function. The instructions may be implemented using any suitable high-level, low-level, object-oriented, visual, compiled and/or interpreted programming language.

(89) One or more aspects of at least one example may be implemented by representative instructions stored on at least one machine-readable medium which represents various logic within the processor, which when read by a machine, computing device or system causes the machine, computing device or system to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

(90) The appearances of the phrase “one example” or “an example” are not necessarily all referring

to the same example or embodiment. Any aspect described herein can be combined with any other aspect or similar aspect described herein, regardless of whether the aspects are described with respect to the same figure or element. Division, omission or inclusion of block functions depicted in the accompanying figures does not infer that the hardware components, circuits, software and/or elements for implementing these functions would necessarily be divided, omitted, or included in embodiments.

(91) Some examples may be described using the expression “coupled” and “connected” along with their derivatives. These terms are not necessarily intended as synonyms for each other. For example, descriptions using the terms “connected” and/or “coupled” may indicate that two or more elements are in direct physical or electrical contact with each other. The term “coupled,” however, may also mean that two or more elements are not in direct contact with each other, but yet still co-operate or interact with each other.

(92) The terms “first,” “second,” and the like, herein do not denote any order, quantity, or importance, but rather are used to distinguish one element from another. The terms “a” and “an” herein do not denote a limitation of quantity, but rather denote the presence of at least one of the referenced items. The term “asserted” used herein with reference to a signal denote a state of the signal, in which the signal is active, and which can be achieved by applying any logic level either logic 0 or logic 1 to the signal. The terms “follow” or “after” can refer to immediately following or following after some other event or events. Other sequences of operations may also be performed according to alternative embodiments. Furthermore, additional operations may be added or removed depending on the particular applications. Any combination of changes can be used and one of ordinary skill in the art with the benefit of this disclosure would understand the many variations, modifications, and alternative embodiments thereof.

(93) Disjunctive language such as the phrase “at least one of X, Y, or Z,” unless specifically stated otherwise, is otherwise understood within the context as used in general to present that an item, term, etc., may be either X, Y, or Z, or any combination thereof (e.g., X, Y, and/or Z). Thus, such disjunctive language is not generally intended to, and should not, imply that certain embodiments require at least one of X, at least one of Y, or at least one of Z to each be present. Additionally, conjunctive language such as the phrase “at least one of X, Y, and Z,” unless specifically stated otherwise, should also be understood to mean X, Y, Z, or any combination thereof, including “X, Y, and/or Z.”

(94) Illustrative examples of the devices, systems, and methods disclosed herein are provided below. An embodiment of the devices, systems, and methods may include any one or more, and any combination of, the examples described below.

(95) References to standards herein can refer to earlier versions, later versions, or derivatives one or more versions of a standard.

(96) Example 1 includes one or more examples, and includes an apparatus comprising: a first graphics processing unit (GPU) comprising at least one integrated network interface controller (NIC), wherein a data plane of the at least one integrated NIC is configured by a particular entity.

(97) Example 2 includes one or more examples, wherein the data plane of the at least one integrated NIC is configured by the particular entity as a control plane but the first GPU and a central processing unit (CPU) are prevented from configuring the data plane of the at least one integrated NIC.

(98) Example 3 includes one or more examples, wherein the data plane of the at least one integrated NIC is configured by the particular entity in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.

(99) Example 4 includes one or more examples, wherein the first GPU is to send packets to a second GPU using the at least one integrated NIC or coupled NIC.

(100) Example 5 includes one or more examples, comprising circuitry to permit the particular

entity to configure the data plane of the at least one integrated NIC but prevent the first GPU or a central processing unit (CPU) to configure the data plane of the at least one integrated NIC.

(101) Example 6 includes one or more examples, comprising a network interface device, wherein the particular entity is to configure the data plane of the at least one integrated NIC using the network interface device.

(102) Example 7 includes one or more examples, comprising a first memory associated with the first GPU and a second memory associated with a second GPU, wherein the first memory comprises a source of data and the second memory comprises a destination of data.

(103) Example 8 includes one or more examples, wherein the first GPU comprises one or more multiprocessors comprising one or more of: single-instruction multiple thread (SIMT) execution circuitry, Single Instruction/Multiple Data (SIMD) execution circuitry, one or more tensor cores, a thread dispatcher, or a thread scheduler.

(104) Example 9 includes one or more examples, comprising: a first server comprising the first GPU.

(105) Example 10 includes one or more examples, comprising: a data center comprising a second server comprising a second GPU communicatively coupled to the first server.

(106) Example 11 includes one or more examples, and includes a non-transitory computer-readable medium comprising instructions stored thereon, that if executed, cause one or more processors to: configure data planes of one or more network interface devices integrated into a graphics processing unit (GPU) by a second network interface device, wherein the second network interface device has exclusive capability to configure the data planes of one or more network interface devices integrated into the GPU.

(107) Example 12 includes one or more examples, wherein a cloud service provider (CSP) is to configure the data planes of one or more network interface devices integrated into the GPU.

(108) Example 13 includes one or more examples, wherein the data planes of one or more network interface devices are to be configured in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.

(109) Example 14 includes one or more examples, wherein the graphics processing unit is to send packets to a second graphics processing unit using the one or more network interface devices.

(110) Example 15 includes one or more examples, wherein a smartNIC is to configure the data planes of the one or more network interface devices but the GPU and a central processing unit (CPU) are not permitted to configure the data planes of the one or more network interface devices.

(111) Example 16 includes one or more examples, and includes a method comprising: receiving a configuration for one or more network interface devices integrated into a graphics processing unit (GPU), wherein the receiving a configuration for one or more network interface devices integrated into the GPU comprises receiving the configuration at a second network interface device and configuring data plane circuitry of the one or more network interface devices based on the configuration, wherein the second network interface device is provided with exclusive capability to configure the data plane circuitry of the network interface devices.

(112) Example 17 includes one or more examples, wherein the data plane circuitry of the one or more network interface devices are to be configured in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.

(113) Example 18 includes one or more examples, wherein the GPU is to send packets to a second graphics processing unit using the one or more network interface devices.

(114) Example 19 includes one or more examples, wherein the second network interface device comprises a smartNIC.

(115) Example 20 includes one or more examples, and includes accessing a memory of a GPU or network connected memory via at least one of the one or more network interface devices.

Claims

1. An apparatus comprising: a first graphics processing unit (GPU) comprising at least one integrated network interface controller (NIC), wherein a data plane of the at least one integrated NIC is configured by a control plane of a second NIC but not by the first GPU and wherein the at least one integrated NIC comprises packet processing circuitry that processes packets and is configured as the data plane.
2. The apparatus of claim 1, wherein the data plane of the at least one integrated NIC is configured by the second NIC as a control plane but the first GPU and a central processing unit (CPU) are prevented from configuring the data plane of the at least one integrated NIC.
3. The apparatus of claim 1, wherein the data plane of the at least one integrated NIC is configured by the second NIC in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.
4. The apparatus of claim 1, wherein the first GPU is to send packets to a second GPU using the at least one integrated NIC.
5. The apparatus of claim 1, comprising circuitry to permit the second NIC to configure the data plane of the at least one integrated NIC but prevent the first GPU or a central processing unit (CPU) to configure the data plane of the at least one integrated NIC.
6. The apparatus of claim 1, comprising a network interface device, wherein the second NIC is to configure the data plane of the at least one integrated NIC using the network interface device.
7. The apparatus of claim 1, comprising a first memory associated with the first GPU and a second memory associated with a second GPU, wherein the first memory comprises a source of data and wherein the first GPU is to cause copying of the data to the second memory.
8. The apparatus of claim 1, wherein the first GPU comprises one or more multiprocessors comprising one or more of: single-instruction multiple thread (SIMT) execution circuitry, Single Instruction/Multiple Data (SIMD) execution circuitry, one or more tensor cores, a thread dispatcher, or a thread scheduler.
9. The apparatus of claim 1, comprising: a first server comprising the first GPU.
10. The apparatus of claim 9, comprising: a data center comprising a second server comprising a second GPU communicatively coupled to the first server.
11. At least one non-transitory computer-readable medium comprising instructions stored thereon, that if executed, cause one or more processors to: configure data planes of one or more network interface devices integrated into a graphics processing unit (GPU) by a control plane of a second network interface device, wherein the second network interface device has exclusive capability to configure data planes of one or more network interface devices integrated into the GPU, wherein the GPU is not to configure the data planes of the one or more network interface devices integrated into the GPU, and wherein the data planes of the one or more network interface devices comprise packet processing circuitries that process packets.
12. The computer-readable medium of claim 11, wherein a cloud service provider (CSP) is to configure the data planes of the one or more network interface devices integrated into the GPU via the second network interface device.
13. The computer-readable medium of claim 11, wherein the data planes of the one or more network interface devices are to be configured in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.
14. The computer-readable medium of claim 11, wherein the graphics processing unit is to send packets to a second graphics processing unit using the one or more network interface devices.
15. The computer-readable medium of claim 11, wherein a central processing unit (CPU) is not

permitted to configure the data planes of the one or more network interface devices.

16. A method comprising: receiving a configuration for one or more network interface devices integrated into a graphics processing unit (GPU), wherein the receiving a configuration for one or more network interface devices integrated into the GPU comprises receiving the configuration from a control plane of a second network interface device and configuring data plane circuitry of the one or more network interface devices based on the configuration, wherein the second network interface device is provided with exclusive capability to configure the data plane circuitries of the network interface devices, wherein the GPU is not permitted to configure the data plane circuitries of the network interface devices, and wherein the data plane circuitries of the one or more network interface devices comprise packet processing circuitries that process packets.

17. The method of claim 16, wherein the data plane circuitries of the one or more network interface devices are to be configured in one or more of: reliability protocol, packet encryption, packet decryption, network protocol to utilize, frame format to utilize, permitted destination for packets, or permitted source for packets.

18. The method of claim 16, wherein the GPU is to send packets to a second graphics processing unit using the one or more network interface devices.

19. The method of claim 16, wherein the second network interface device comprises a smartNIC.

20. The method of claim 16, comprising: accessing a memory of a GPU or network connected memory via at least one of the one or more network interface devices.
