

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12413530
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Shenbagam; Raghunath et al.

Data processing system with link-based resource allocation for reconfigurable processors

Abstract

The disclosed technology relates to link-based resource allocation for a pool of reconfigurable processors. Resource allocation is provided for reconfigurable processors based on link bandwidths and link latencies. Runtime logic receives target link bandwidth and target link latency and rated link bandwidth and rated link latency. In response, the runtime logic allocates configuration files for an application, reconfigurable processors, and links between the processors. The runtime logic executes the allocated configuration files using the allocated processors and the allocated links. In another embodiment, the pool of reconfigurable processors comprise a cluster of processing nodes connected through a network.

Inventors:	Shenbagam; Raghunath (San Jose, CA), Kumar; Ravinder (Fremont, CA)
Applicant:	SambaNova Systems, Inc. (Palo Alto, CA)
Family ID:	78828598
Assignee:	SambaNova Systems, Inc. (Palo Alto, CA)
Appl. No.:	18/740240
Filed:	June 11, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240330074 A1	Oct. 03, 2024

Related U.S. Application Data

continuation parent-doc US 17528081 20211116 US 12008417 20240611 child-doc US 18740240
continuation parent-doc US 17214768 20210326 US 11200096 20211214 child-doc US 17528081

Publication Classification

Int. Cl.: **H04L47/28** (20220101); **G06F9/445** (20180101); **G06F9/455** (20180101); **G06F9/50** (20060101); **G06F15/78** (20060101); **H04L41/0895** (20220101); **H04L41/0896** (20220101); **H04L45/745** (20220101); **H04L49/10** (20220101)

U.S. Cl.:

CPC **H04L47/28** (20130101); **G06F9/44505** (20130101); **G06F9/45558** (20130101); **G06F9/5044** (20130101); **G06F9/5077** (20130101); **G06F15/7892** (20130101); **H04L41/0896** (20130101); **H04L49/10** (20130101); G06F2009/4557 (20130101); G06F2009/45595 (20130101); G06F2209/5011 (20130101); G06F2213/0062 (20130101); G06F2213/0064 (20130101); H04L41/0895 (20220501); H04L45/745 (20130101)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
5684980	12/1996	Casselman	N/A	N/A
6470485	12/2001	Cote et al.	N/A	N/A
6539438	12/2002	Ledzius et al.	N/A	N/A
6667983	12/2002	Lo et al.	N/A	N/A
8745626	12/2013	Sandstrom	712/15	G06F 15/7867
9009723	12/2014	Degenaro et al.	N/A	N/A
9501325	12/2015	Pell et al.	N/A	N/A
10511479	12/2018	Xie et al.	N/A	N/A
10621138	12/2019	Hu et al.	N/A	N/A
10698853	12/2019	Grohoski et al.	N/A	N/A
10768899	12/2019	Koeplinger et al.	N/A	N/A
10802870	12/2019	Lu	N/A	N/A
10831507	12/2019	Shah et al.	N/A	N/A
10831523	12/2019	Kochevar-Cureton et al.	N/A	N/A
10877822	12/2019	Wang et al.	N/A	N/A
11055141	12/2020	Prabhakar et al.	N/A	N/A
11068780	12/2020	Mellempudi et al.	N/A	N/A
11080227	12/2020	Koeplinger et al.	N/A	N/A
11126574	12/2020	Prabhakar et al.	N/A	N/A
11150872	12/2020	Wang et al.	N/A	N/A
11182221	12/2020	Sivaramakrishnan et al.	N/A	N/A
11182264	12/2020	Sivaramakrishnan et al.	N/A	N/A

11184439	12/2020	Eran et al.	N/A	N/A
11188497	12/2020	Shah et al.	N/A	N/A
11200096	12/2020	Shenbagam et al.	N/A	N/A
11237880	12/2021	Raumann et al.	N/A	N/A
11237971	12/2021	Brown et al.	N/A	N/A
11250105	12/2021	Wang et al.	N/A	N/A
11327713	12/2021	Wang et al.	N/A	N/A
11327717	12/2021	Wang et al.	N/A	N/A
11327923	12/2021	Wang et al.	N/A	N/A
11328038	12/2021	Wang et al.	N/A	N/A
11328207	12/2021	Lauterbach et al.	N/A	N/A
11328208	12/2021	Lie et al.	N/A	N/A
11347965	12/2021	Dutta et al.	N/A	N/A
11360800	12/2021	Kochevar-Cureton et al.	N/A	N/A
11386038	12/2021	Prabhakar et al.	N/A	N/A
11392740	12/2021	Raumann et al.	N/A	N/A
11410027	12/2021	Chen et al.	N/A	N/A
11436429	12/2021	Jaganathan et al.	N/A	N/A
11645057	12/2022	Koeplinger et al.	N/A	N/A
11709664	12/2022	Chen et al.	N/A	N/A
11782729	12/2022	Grohoski et al.	N/A	N/A
11809908	12/2022	Kumar et al.	N/A	N/A
11836629	12/2022	Liu	N/A	N/A
11886930	12/2023	Sivaramakrishnan et al.	N/A	N/A
2002/0156998	12/2001	Casselman	N/A	N/A
2003/0108119	12/2002	Mohebbi et al.	N/A	N/A
2006/0012395	12/2005	Huppenthal et al.	N/A	N/A
2006/0015712	12/2005	Ang et al.	N/A	N/A
2007/0186126	12/2006	Smith et al.	N/A	N/A
2007/0220522	12/2006	Coene et al.	N/A	N/A
2008/0013448	12/2007	Horie et al.	N/A	N/A
2009/0089475	12/2008	Chitlur	N/A	N/A
2009/0172351	12/2008	Vorbach et al.	N/A	N/A
2009/0300209	12/2008	Elzur	N/A	N/A
2014/0137123	12/2013	Hartmann et al.	N/A	N/A
2014/0258438	12/2013	Ayoub et al.	N/A	N/A
2015/0058614	12/2014	Degenaro et al.	N/A	N/A
2015/0100971	12/2014	Dube et al.	N/A	N/A
2015/0106823	12/2014	Canoy et al.	N/A	N/A
2016/0308719	12/2015	Putnam et al.	N/A	N/A
2016/0314025	12/2015	Mcgarry et al.	N/A	N/A
2016/0378550	12/2015	Monfort et al.	N/A	N/A
2017/0220499	12/2016	Gray	N/A	N/A
2017/0289060	12/2016	Aftab	N/A	H04L 67/1097
2017/0315815	12/2016	Smith et al.	N/A	N/A
2017/0317679	12/2016	Suh et al.	N/A	N/A
2018/0285295	12/2017	Abel et al.	N/A	N/A

2018/0307950	12/2017	Nealis et al.	N/A	N/A
2018/0308200	12/2017	Surti et al.	N/A	N/A
2018/0314941	12/2017	Lie et al.	N/A	N/A
2018/0315158	12/2017	Nurvitadhi et al.	N/A	N/A
2019/0089616	12/2018	Chabbi et al.	N/A	N/A
2019/0138890	12/2018	Liang et al.	N/A	N/A
2019/0171604	12/2018	Brewer	N/A	N/A
2019/0171612	12/2018	Shahar et al.	N/A	N/A
2019/0180176	12/2018	Yudanov et al.	N/A	N/A
2019/0258921	12/2018	Lie et al.	N/A	N/A
2019/0286973	12/2018	Kovvuri et al.	N/A	N/A
2019/0347136	12/2018	Miyoshi	N/A	G06F 11/3031
2019/0384642	12/2018	Bolkhovitin et al.	N/A	N/A
2020/0090313	12/2019	Bugdary et al.	N/A	N/A
2020/0142753	12/2019	Harwood et al.	N/A	N/A
2020/0142857	12/2019	Catiller et al.	N/A	N/A
2020/0151573	12/2019	Das et al.	N/A	N/A
2020/0174840	12/2019	Zhao et al.	N/A	N/A
2020/0183745	12/2019	Ernst et al.	N/A	N/A
2020/0226444	12/2019	Sharma et al.	N/A	N/A
2020/0264876	12/2019	Lo et al.	N/A	N/A
2020/0301898	12/2019	Samynathan et al.	N/A	N/A
2020/0314181	12/2019	Eran	N/A	N/A
2020/0326992	12/2019	Jin et al.	N/A	N/A
2020/0341930	12/2019	Cannata et al.	N/A	N/A
2021/0011770	12/2020	Prabhakar et al.	N/A	N/A
2021/0081691	12/2020	Chen et al.	N/A	N/A
2021/0089343	12/2020	Hyoudou	N/A	N/A
2021/0097366	12/2020	Wagner et al.	N/A	N/A
2021/0097379	12/2020	Yang et al.	N/A	N/A
2021/0103820	12/2020	Ghosh	N/A	N/A
2021/0125058	12/2020	Chowdhury et al.	N/A	N/A
2021/0192357	12/2020	Sinha et al.	N/A	N/A
2021/0192358	12/2020	Song et al.	N/A	N/A
2021/0200610	12/2020	Chu et al.	N/A	N/A
2021/0241093	12/2020	Byrne et al.	N/A	N/A
2021/0374503	12/2020	Kim et al.	N/A	N/A
2022/0058034	12/2021	Grohoski et al.	N/A	N/A
2022/0121928	12/2021	Dong	N/A	H04L 41/0893
2022/0197712	12/2021	Sivaramakrishnan et al.	N/A	N/A
2022/0197714	12/2021	Raumann et al.	N/A	N/A
2022/0198117	12/2021	Raumann et al.	N/A	N/A
2022/0269534	12/2021	Misra et al.	N/A	N/A
2022/0308935	12/2021	Shenbagam et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1372084	12/2002	EP	N/A
2020112901	12/2019	JP	N/A
201606508	12/2015	TW	N/A
202240386	12/2021	TW	N/A
202240394	12/2021	TW	N/A
202248853	12/2021	TW	N/A
2010142987	12/2009	WO	N/A
2022133047	12/2021	WO	N/A
2022182573	12/2021	WO	N/A
2022203925	12/2021	WO	N/A

OTHER PUBLICATIONS

Accelerated Computing with a Reconfigurable Dataflow Architecture, SambaNova Systems Whitepaper, 10 pages. cited by applicant

Bae et al., Auto-Tuning CNNs for Coarse-Grained Reconfigurable Array-based Accelerators, IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems, vol. 37, Issue: 11, Nov. 2018, 10 pages. cited by applicant

Busa et al., A Run-Time Word-Level Reconfigurable Coarse-Grain Functional Unit for a VLIW processor; ACM 2002, pp. 44-49. (Year: 2002). cited by applicant

Dettmers, How to Parallelize Deep Learning on GPUs Part 1 of 2: Data Parallelism, dated Oct. 9, 2014, 19 pages. Retrieved on Sep. 3, 2021 Retrieved from [URL:

<https://timdettmers.com/2014/10/09/deep-learning-data-parallelism/>]. cited by applicant

Dettmers, How to Parallelize Deep Learning on GPUs Part 2 of 2: Model Parallelism, dated Nov. 9, 2014, 19 pages. Retrieved on Sep. 3, 2021. Retrieved from [URL:

<https://timdettmers.com/2014/11/09/model-parallelism-deep-learning/>]. cited by applicant

Donges, Gradient Descent: An Introduction to Machine Learning's Most Popular Algorithms, dated Jun. 16, 2019, 10 pages. Retrieved on Mar. 24, 2021, retrieved from [URL: <https://builtin.com/data-science/gradient-descent>]. cited by applicant

Ekanayake, Model Parallelism in Deep Learning is NOT What you think, dated Nov. 10, 2018, 4 pages. Retrieved an Sep. 3, 2021. Retrieved from [URL: <https://medium.com/@esaliya/model-parallelism-in-deep-learning-is-not-what-you-think-94d2f81e82ed>]. cited by applicant

Éricles Sousa, A reconfigurable memory architecture for system integration of coarse-grained reconfigurable arrays, Published in: 2017 International Conference on ReConFigurable Computing and FPGAs (ReConFig) Dec. 4-6, 2017 <https://ieeexplore.IEEE.org/document/8279768> (Year: 2017). cited by applicant

Fazlali et al., Efficient task scheduling for runtime reconfigurable systems, Journal of Systems Architecture, vol. 56, dated Jul. 26, 2010, pp. 623-632, 10 pages. cited by applicant

Galanis et al., A design flow for speeding-up dsp applications in heterogeneous reconfigurable systems, Microelectronics Journal, vol. 37, dated 2006, pp. 554-564, 11 pages. cited by applicant

Galanis et al., Accelerating Applications by Mapping Critical Kernels on Coarse-Grain Reconfigurable Hardware in Hybrid Systems, Field-Programmable Custom Computing Machines, 2005, 13th Annual IEEE Symposium on Napa, CA Apr. 2005, 2 pages. cited by applicant

Galanis et al., Partitioning Methodology for Heterogeneous Reconfigurable Functional Units, The Journal of Supercomputing, vol. 38, No. 1, dated Oct. 1, 2006, 18 pages. cited by applicant

Goodfellow et. al., Deep Learning Book Chapter 6 Deep Feedforward Networks, 2016, 60 pages. cited by applicant

Insujang, GPU Architecture Overview, Better Tomorrow with Computer Science, published Apr. 27, 2017, retrieved on Jun. 17, 2021, retrieved from the Internet [URL:

<https://insujang.github.io/2017-04-17/gpu-architecture-overview/>]. cited by applicant

Iqbal et al., Reconfigurable Processor Architecture for High Speed Applications, IEEE, dated 2009, pp. 624-629. cited by applicant

Jackson et al., PCI Express Technology Comprehensive Guide to Generation 1.x, 2.x and 3.0, dated Jun. 2020, 1057 pages. cited by applicant

Jafri et al., NeuroCGRA: A CGRAs with Support for Neural Networks, 2014 International Conference on High Performance Computing & Simulation (HPCS), 8 pages. cited by applicant

Jin et. al., How to scale distributed deep learning, dated Nov. 14, 2016, 16 pages. cited by applicant

Kachris et al.; "A Survey on Reconfigurable Accelerators for Cloud Computing", IEEE 2016, Aug. 29, 2016, pp. 1-11. cited by applicant

Knodel, Oliver, et. al., "RC3E: Reconfigurable Accelerators in Data Centers and their Provision by Adapted Service Models", IEEE 9th International Convergence on Cloud Computing, 2016, pp. 1-8. cited by applicant

Koeplinger et al., Spatial: A Language and Compiler for Application Accelerators, PLDI '18, Jun. 18-22, 2018, Association for Computing Machinery, 16 pages. cited by applicant

Lecture 11: Distributed Training and Communication Protocols, CSE599W: Spring 2018, UW Paul G. Allen School of Computer Science and Engineering, 41 pages. cited by applicant

Li, Ang, et. al., "Evaluating Modern GPU Interconnect: PCIe, NVLink, NV-SLI, NVSwitch and GPUDirect", Mar. 11, 2019, 15 pages. cited by applicant

Li, et al., "CATERPILLAR: Coarse Grain Reconfigurable Architecture for Accelerating the Training of Deep Neural Networks," arXiv: 1706.00517v2 [cs.DC], Jun. 8, 2017, 10 pages. cited by applicant

Liang et al., Dynamic Coarse Grain Dataflow Reconfiguration Technique for Real-Time Systems, IEEE, dated 2005, pp. 3511-3514. cited by applicant

Liu et. al., Offloading distributed Applications onto SmartNICs using iPipe, ACM 2019, pp. 1-16. cited by applicant

M. Emani et al., Accelerating Scientific Applications With Sambanova Reconfigurable Dataflow Architecture, in Computing in Science & Engineering, vol. 23, No. 2, pp. 114-119, Mar. 26, 2021, [doi: 10.1109/MCSE.2021.3057203]. cited by applicant

Ma et al., DeepGauge: Multi-Granularity Testing Criteria for Deep Learning Systems; ACM 2018, pp. 1-12. cited by applicant

Mao, Data Parallelism vs Model Parallelism in Distributed Deep Learning Training, dated Mar. 23, 2019, 4 pages, retrieved on Mar. 30, 2021, Retrieved from the internet [URL: <https://leimao.github.io>]. cited by applicant

Marshall, Dave, "Remote Procedure Calls (RPC)", Jan. 5, 1999, 15 pages, Retrieved from URL. cited by applicant

Mazur, A step by step backpropagation example, dated Mar. 17, 2015, 26 pages. Retrieved on Sep. 3, 2021. Retrieved from [URL: <https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>]. cited by applicant

NVIDIA, "NVIDIA DGX-1 System Architecture", WP-08437-001_v02, 2017, 33 pages. cited by applicant

NVIDIA, "NVIDIA DGX-1 With Tesla V100 System Architecture", WP-08437-002_v01, 2017, 43 pages. cited by applicant

NVIDIA, "NVIDIA Tesla P100", WP-08019-001 v01.1, 2016, 45 pages. cited by applicant

NVIDIA, "NVIDIA Turing GPU Architecture", WP-09183-001_v01, 2018, 86 pages. cited by applicant

Padole et al., Configuration Memory Based Dynamic Coarse Grained Reconfigurable Multiscore Architecture, IEEE 2013, pp. 3511-3514, 5 pages. cited by applicant

Paek et al., "Binary Acceleration Using Coarse-Grained Reconfigurable Architecture," ACM SIGARCH Computer Architecture News, vol. 38, No. 4, Sep. 2010, 7 pages. cited by applicant

PCT/US2021/063728—International Search Report and Written Opinion, dated Apr. 4, 2022, 15 pages. cited by applicant

PCT/US2021/063733—International Search Report and Written Opinion, dated Apr. 4, 2022, 17 pages. cited by applicant

PCT/US2022/016871—International Search Report and Written Opinion, dated Jun. 1, 2022, 14 pages. cited by applicant

PCT/US2022/020638—International Search Report and Written Opinion, dated Jun. 21, 2022, 17 pages. cited by applicant

Podobas et al, A Survey on Coarse-Grained Reconfigurable Architectures From a Performance Perspective, IEEEAccess, vol. 2020.3012084, Jul. 27, 2020, 25 pages. cited by applicant

Prabhakar et al., Plasticine: A Reconfigurable Architecture for Parallel Patterns, ISCA, Jun. 24-28, 2017, 14 pages. cited by applicant

Rubattu et al., Dataflow-Functional High-Level Synthesis for Coarse-Grained Reconfigurable Accelerators, IEEE 2019, pp. 69-72 (Year: 2019). cited by applicant

Ruder, An overview of gradient descent optimization algorithms, NUI Galway Aylien Lyd, dated Jun. 15, 2017, 14 pages. cited by applicant

Souissi et al., Optimization of Run-time Mapping on Heterogeneous CPU/FPGA Architecture, 9th International Conference of Modeling, Optimization and Simulation—MOSIM'12, Jun. 6-8, 2012, Bordeaux, France, 9 pages. cited by applicant

Strom, Scalable Distributed DNN Training Using Commodity GPU Cloud Computing, Amazon.com, 5 pages. cited by applicant

Tanaka et. al., Distributed Deep Learning with GPU-FPGA heterogenous computing, IEEE 2021, 9 pages. cited by applicant

TW 110147198, Allowance Decision, Jul. 27, 2022, 2 pages. cited by applicant

TW 110147198, Search Report, Jul. 27, 2022, 1 page. cited by applicant

TW 111110601, Non-final office action, dated Oct. 18, 2022, 3 pages. cited by applicant

TW 111110601, Search Report, Oct. 18, 2022, 1 page. cited by applicant

U.S. Appl. No. 17/127,818 Notice of Allowance, dated Jul. 21, 2021, 10 pages. cited by applicant

U.S. Appl. No. 17/127,818—Office Action dated Apr. 1, 2021, 15 pages. cited by applicant

U.S. Appl. No. 17/127,929 Notice of Allowance, dated Jul. 21, 2021, 14 pages. cited by applicant

U.S. Appl. No. 17/127,929—Office Action dated Apr. 1, 2021, 26 pages. cited by applicant

U.S. Appl. No. 17/127,818—Response to Office Action dated Apr. 1, 2021, filed Jul. 1, 2021, 15 pages. cited by applicant

U.S. Appl. No. 17/127,929—Response to Office Action dated Apr. 1, 2021, filed Jul. 1, 2021, 10 pages. cited by applicant

U.S. Appl. No. 17/185,264—Non-Final Office Action, dated Jan. 26, 2023, 9 pages. cited by applicant

U.S. Appl. No. 17/185,264—Notice of Allowance, dated May 30, 2023, 7 pages. cited by applicant

U.S. Appl. No. 17/214,768—Notice of Allowance, dated Aug. 11, 2021, 26 pages. cited by applicant

U.S. Appl. No. 17/214,768—Supplemental Notice of Allowance, dated Aug. 25, 2021, 10 pages. cited by applicant

U.S. Appl. No. 17/379,921—Notice of Allowance, dated Nov. 26, 2021, 21 pages. cited by applicant

U.S. Appl. No. 17/379,921 Notice of Allowance dated Mar. 21, 2022, 25 pages. cited by applicant

U.S. Appl. No. 17/379,924—Notice of Allowance, dated Sep. 16, 2021, 23 pages. cited by applicant

U.S. Appl. No. 17/522,655—Notice of Allowance, dated Nov. 16, 2022, 21 pages. cited by applicant

U.S. Appl. No. 17/522,658—Notice of Allowance, dated Dec. 14, 2022, 22 pages. cited by applicant

applicant
U.S. Appl. No. 17/522,682—Notice of Allowance, dated Jan. 11, 2023, 23 pages. cited by applicant
U.S. Appl. No. 17/522,682—Supplemental Notice of Allowance, dated Jan. 24, 2023, 2 pages.
cited by applicant
U.S. Appl. No. 17/522,694—Non-Final Office Action, dated Mar. 31, 2023, 29 pages. cited by
applicant
U.S. Appl. No. 17/522,694—Notice of Allowance, dated Sep. 6, 2023, 62 pages. cited by applicant
U.S. Appl. No. 17/528,081—Non-Final Office Action, dated Jun. 7, 2023, 29 pages. cited by
applicant
Vucha et al., Dynamic Task Distribution Model for On-Chip Reconfigurable High Speed
Computing System, Hindawi, dated Jun. 30, 2015, 13 pages. cited by applicant
What is the difference between model paralellism and data paralellism, Quora, 27 pages. Retrieved
on Sep. 3, 2021. Retrieved from [URL: <https://www.quora.com/What-is-the-difference-between-model-parallelism-and-data-parallelism>]. cited by applicant
Woolloy, NCCL: Accelerated Multi-GPU Collective Communications, NVIDIA, 56 pages. cited by
applicant
Xiandong Qi, Introduction to Distributed Deep Learning, dated May 13, 2017, 13 pages. cited by
applicant
Zhang et. al., Dive into Deep Learning, Release 0.16.2, dated Mar. 20, 2021, 1027 pages. cited by
applicant

Primary Examiner: Teets; Bradley A

Assistant Examiner: Headly; Melissa A

Attorney, Agent or Firm: Flagship Patents

Background/Summary

PRIORITY APPLICATION (1) This application is a continuation of U.S. Nonprovisional patent application Ser. No. 17/528,081, filed Nov. 16, 2021, entitled “Interconnect-Based Resource Allocation For Reconfigurable Processors”, now U.S. Pat. No. 12,008,417, issued Jun. 11, 2024, which is a continuation of U.S. Nonprovisional patent application Ser. No. 17/214,768, filed Mar. 26, 2021, entitled “Resource Allocation for Reconfigurable Processors”, now U.S. Pat. No. 11,200,096, issued Dec. 14, 2021, both of which are incorporated herein by reference for all purposes.

INCORPORATIONS

(1) The following are incorporated by reference for all purposes as if fully set forth herein: Prabhakar et al., “Plasticine: A Reconfigurable Architecture for Parallel Patterns,” ISCA '17, Jun. 24-28, 2017, Toronto, ON, Canada; Koeplinger et al., “Spatial: A Language and Compiler for Application Accelerators,” *Proceedings Of The 39th ACM SIGPLAN Conference On Programming Language Design And Implementation (PLDI)*, Proceedings of the 43rd International Symposium on Computer Architecture, 2018; U.S. Nonprovisional patent application Ser. No. 16/239,252, filed Jan. 3, 2019, titled, “VIRTUALIZATION OF A RECONFIGURABLE DATA PROCESSOR,” now U.S. Pat. No. 10,698,853, issued on Jun. 30, 2020; U.S. Nonprovisional patent application Ser. No. 16/197,826, filed Nov. 21, 2018, titled, “CONFIGURATION LOAD OF A RECONFIGURABLE DATA PROCESSOR,” now U.S. Pat. No. 10,831,507, issued on Nov. 10, 2020; U.S. Nonprovisional patent application Ser. No. 16/198,086, filed Nov. 21, 2018, titled,

“CONFIGURATION UNLOAD OF A RECONFIGURABLE DATA PROCESSOR,” now U.S. Pat. No. 11,188,497, issued on Nov. 30, 2021; U.S. Nonprovisional patent application Ser. No. 16/260,548, filed Jan. 29, 2019, titled, “MATRIX NORMAL/TRANPOSE READ AND A RECONFIGURABLE DATA PROCESSOR INCLUDING SAME,” now U.S. Pat. No. 10,768,899, issued on Sep. 8, 2020; U.S. Nonprovisional patent application Ser. No. 16/536,192, filed Aug. 8, 2019, titled, “COMPILER FLOW LOGIC FOR RECONFIGURABLE ARCHITECTURES,” now U.S. Pat. No. 11,080,227, issued on Aug. 3, 2021; U.S. Nonprovisional patent application Ser. No. 16/407,675, filed May 9, 2019, titled, “CONTROL FLOW BARRIER AND RECONFIGURABLE DATA PROCESSOR,” now U.S. Pat. No. 11,386,038, issued on Jul. 12, 2022; U.S. Nonprovisional patent application Ser. No. 16/504,627, filed Jul. 8, 2019, titled, “QUIESCE RECONFIGURABLE DATA PROCESSOR,” now U.S. Pat. No. 11,055,141, issued on Jul. 6, 2021; U.S. Nonprovisional patent application Ser. No. 16/572,516, filed Sep. 16, 2019, titled, “EFFICIENT EXECUTION OF OPERATION UNIT GRAPHS ON RECONFIGURABLE ARCHITECTURES BASED ON USER SPECIFICATION”; U.S. Nonprovisional patent application Ser. No. 16/744,077, filed Jan. 15, 2020, titled, “COMPUTATIONALLY EFFICIENT SOFTMAX LOSS GRADIENT BACKPROPAGATION,” now U.S. Pat. No. 11,836,629, issued on Dec. 5, 2023; U.S. Nonprovisional patent application Ser. No. 16/590,058, filed Oct. 1, 2019, titled, “COMPUTATION UNITS FOR FUNCTIONS BASED ON LOOKUP TABLES,” now U.S. Pat. No. 11,327,713, issued on May 10, 2022; U.S. Nonprovisional patent application Ser. No. 16/695,138, filed Nov. 25, 2019, titled, “COMPUTATIONAL UNITS FOR BATCH NORMALIZATION,” now U.S. Pat. No. 11,328,038, issued on May 10, 2022; U.S. Nonprovisional patent application Ser. No. 16/688,069, filed Nov. 19, 2019, titled, “LOOK-UP TABLE WITH INPUT OFFSETTING,” now U.S. Pat. No. 11,327,717, issued on May 10, 2022; U.S. Nonprovisional patent application Ser. No. 16/718,094, filed Dec. 17, 2019, titled, “COMPUTATIONAL UNITS FOR ELEMENT APPROXIMATION,” now U.S. Pat. No. 11,150,872, issued on Oct. 19, 2021; U.S. Nonprovisional patent application Ser. No. 16/560,057, filed Sep. 4, 2019, titled, “SIGMOID FUNCTION IN HARDWARE AND A RECONFIGURABLE DATA PROCESSOR INCLUDING SAME,” now U.S. Pat. No. 11,327,923, issued on May 10, 2022; U.S. Nonprovisional patent application Ser. No. 16/572,527, filed Sep. 16, 2019, titled, “PERFORMANCE ESTIMATION-BASED RESOURCE ALLOCATION FOR RECONFIGURABLE ARCHITECTURES,” now U.S. Pat. No. 11,410,027, issued on Aug. 9, 2022; U.S. Nonprovisional patent application Ser. No. 15/930,381, filed May 12, 2020, titled, “COMPUTATIONALLY EFFICIENT GENERAL MATRIX-MATRIX MULTIPLICATION (GeMM),” now U.S. Pat. No. 11,250,105, issued on Feb. 15, 2022; U.S. Nonprovisional patent application Ser. No. 16/890,841, filed Jun. 2, 2020, titled, “ANTI-CONGESTION FLOW CONTROL FOR RECONFIGURABLE PROCESSORS,” now U.S. Pat. No. 11,709,664, issued on Jul. 25, 2023; U.S. Nonprovisional patent application Ser. No. 17/023,015, filed Sep. 16, 2020, titled, “COMPILE TIME LOGIC FOR DETECTING STREAMING COMPATIBLE AND BROADCAST COMPATIBLE DATA ACCESS PATTERNS,” now U.S. Pat. No. 11,237,971, issued on Feb. 1, 2022; U.S. Nonprovisional patent application Ser. No. 17/031,679, filed Sep. 24, 2020, titled, “SYSTEMS AND METHODS FOR MEMORY LAYOUT DETERMINATION AND CONFLICT RESOLUTION,” now U.S. Pat. No. 11,645,057, issued on May 9, 2023; U.S. Nonprovisional patent application Ser. No. 17/175,289, filed Feb. 12, 2021, titled, “INSTRUMENTATION PROFILING FOR RECONFIGURABLE PROCESSORS,” now U.S. Pat. No. 11,126,574, issued on Sep. 21, 2021; U.S. Nonprovisional patent application Ser. No. 16/922,975, filed Jul. 7, 2020, titled, “RUNTIME VIRTUALIZATION OF RECONFIGURABLE DATAFLOW RESOURCES,” now U.S. Pat. No. 11,809,908, issued on Nov. 7, 2023; U.S. Nonprovisional patent application Ser. No. 16/996,666, filed Aug. 18, 2020, titled, “RUNTIME PATCHING OF CONFIGURATION FILES,” now U.S. Pat. No. 11,782,729, issued on Oct. 10, 2023; U.S. Nonprovisional patent application Ser. No. 17/127,818, filed Dec. 18, 2020, titled,

“INTRA-NODE BUFFER-BASED STREAMING FOR RECONFIGURABLE PROCESSOR-AS-A-SERVICE (RPaaS),” now U.S. Pat. No. 11,182,264, issued on Nov. 23, 2021; and U.S. Nonprovisional patent application Ser. No. 17/127,929, filed Dec. 18, 2020, titled, “INTER-NODE BUFFER-BASED STREAMING FOR RECONFIGURABLE PROCESSOR-AS-A-SERVICE (RPaaS),” now U.S. Pat. No. 11,182,221, issued on Nov. 23, 2021.

FIELD OF THE TECHNOLOGY DISCLOSED

(2) The technology disclosed relates to bandwidth and latency optimization in systems and methods that use processors like Central Processing Units (CPUs), Graphics Processing Units (GPUs), Field Programmable Gate Arrays (FPGAs), Coarse-Grained Reconfigurable Architectures (CGRAs), Application-Specific Integrated Circuits (ASICs), Application Specific Instruction-set Processor (ASIP), and Digital Signal Processors (DSPs). In particular, the technology disclosed relates to resource allocation for reconfigurable processors based on link bandwidths and link latencies.

BACKGROUND

(3) The subject matter discussed in this section should not be assumed to be prior art merely as a result of its mention in this section. Similarly, a problem mentioned in this section or associated with the subject matter provided as background should not be assumed to have been previously recognized in the prior art. The subject matter in this section merely represents different approaches, which in and of themselves can also correspond to implementations of the claimed technology.

(4) Reconfigurable accelerators, including Field Programmable Gate Arrays (FPGAs), can be configured to implement a variety of functions more efficiently or faster than might be achieved using a general purpose processor executing a computer program. So-called Coarse-Grained Reconfigurable Architectures (CGRAs) are being developed in which configurable units in an array are more complex than used in typical, more fine-grained FPGAs, and can enable faster or more efficient execution of various classes of functions. For example, CGRAs have been proposed that can enable implementation of energy-efficient accelerators for machine learning and artificial intelligence workloads. See, Prabhakar, et al., “Plasticine: A Reconfigurable Architecture for Parallel Patterns,” ISCA '17, Jun. 24-28, 2017, Toronto, ON, Canada.

(5) Deep neural networks learn many levels of abstraction, ranging from simple to complex concepts. The strength of deep models is that they are not only powerful but learnable. A deep neural network is trained by feeding it input and letting it compute layer-by-layer to generate output for comparison with a known correct answer. After computing the error at the output, this error flows backward through the network by back-propagation. At each step backward the model parameters are tuned in a direction that tries to reduce the error using a numerical optimization method such as stochastic gradient descent (SGD). This process sweeps over the data improving the model as it goes.

(6) To serve increasingly sophisticated applications and deliver higher quality results, neural networks are becoming exponentially deeper and more complex. At the same time, neural networks deployed in mainstream interactive applications are being driven to infer and predict results faster. As neural networks get deeper and more complex, they provide a dramatic increase in accuracy but training these higher accuracy networks requires much higher computation time, and their complexity increases prediction latency.

(7) Servers with two or more accelerators per host are becoming common as developers increasingly expose and leverage the available parallelism in their applications. While dense accelerator systems provide a great vehicle for scaling single-node performance, multi-accelerator application efficiency can be constrained by the performance of the PCIe (Peripheral Component Link Express) bus connections between accelerators. Similarly, data center applications are growing outside the server enclosure, requiring efficient scaling across multiple linked systems.

(8) Training deep neural networks in parallel across multiple accelerators and/or multiple nodes requires distributing either the input data (“data parallel”), the model being trained (“model

parallel”), or a hybrid of the two. Regardless of the approach, parallelizing across accelerators requires synchronization and communication of data (such as gradients) between accelerators. For example, in data-parallel approaches, separate parallel tasks must periodically resynchronize their gradients so that the model parameters are kept in sync across all parallel tasks. This amounts to an all reduce operation.

(9) Scaling is a measure of the improvement in time to solution when increasing the number of parallel processors applied to a problem. A common approach to scaling training deep neural networks is to increase the global batch size as the number of accelerators increases. Perhaps unsurprisingly, these so-called “weak” scaling approaches have high parallel efficiency, even with relatively slow links among accelerators.

(10) A common deep learning workload is the training of convolutional neural networks for computer vision tasks such as image recognition and understanding. Implementations commonly use weak data-parallel scaling for training. The training of recurrent neural networks (RNNs) is another common deep learning workload. Due to high inter-accelerator communication requirements, RNNs—in particular, long short-term memory (LSTM) networks—often benefit more from faster links. RNNs are commonly used for speech recognition and natural language processing. An example is the “Sequence-to-Sequence” (Seq2Seq) neural machine translation (NMT) technique, with several implementations and improvements including Open Neural Machine Translation (Open MT).

(11) Seq2Seq NMT models are used for language translation as well as image and video captioning. Seq2Seq is an RNN made up of an encoder, which ingests an input sequence, and a decoder, which generates the output sequence (e.g., a translation). The encoder and decoder are parameterized by their embedding size, known as the RNN's learning capacity. Increasing the embedding size can improve the network's accuracy at the cost of lower training throughput (as with any other network, care must be taken to not overfit). With larger embeddings, the performance of the links becomes more important.

(12) Multi-system scaling of the latest computational workloads, especially deep learning, requires strong communications between accelerators, both inside the system and between systems, to match the significant accelerator performance of each system and to improve performance of scaling workloads. An opportunity arises to improve bandwidth and reduce latency in accelerator implementations of deep learning models.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) In the drawings, like reference characters generally refer to like parts throughout the different views. Also, the drawings are not necessarily to scale, with an emphasis instead generally being placed upon illustrating the principles of the technology disclosed. In the following description, various implementations of the technology disclosed are described with reference to the following drawings, in which.
- (2) FIG. 1 shows one implementation of an architectural level schematic depicting elements in a first processing node of a data center.
- (3) FIG. 2A shows an example specification table for target link bandwidths of a first processing node in the data center of FIG. 1.
- (4) FIG. 2B shows an example specification table for rated link bandwidths of a first processing node in the data center of FIG. 1.
- (5) FIG. 2C shows an example color-coded schematic comparing target link bandwidths to rated link bandwidths of a first processing node in the data center of FIG. 1.
- (6) FIG. 3 shows an example runtime logic assignment of rated link latencies in a color-coded

schematic as normalized scores from one to one hundred for a first processing node in the data center.

(7) FIG. 4A shows an example table representing a first specification for target link latency as normalized scores of a first processing node in the data center of FIG. 1.

(8) FIG. 4B shows an example table representing a resource allocator routing table of rated links numbered as normalized scores for a first processing node in the data center of FIG. 1.

(9) FIG. 4C shows an example color-coded schematic comparing virtual pairs for matching the normalized scores of target link latency to resource allocator routing table normalized scores of rated links latency in a first processing node in the data center of FIG. 1.

(10) FIG. 5A shows an example color-coded schematic according to a revised resource allocator routing table's matching normalized scores for a first processing node in the data center.

(11) FIG. 5B shows an example table representing a revised resource allocator routing table of normalized scores for a first application in a first processing node.

(12) FIG. 5C shows an example color-coded schematic according to a revised resource allocator routing table's matching normalized scores for a first processing node in the data center.

(13) FIG. 6A shows an example table representing a second specification for a target link latency numbered as normalized scores for a first processing node in the data center.

(14) FIG. 6B, same as FIG. 4B, shows an example table representing a resource allocator routing table of rated links numbered as normalized scores for a first processing node.

(15) FIG. 6C shows an example color-coded schematic comparing virtual pairs for matching the second specification's target link latency as normalized scores to resource allocator routing table's rated links latency as normalized scores for a first processing node.

(16) FIG. 7A shows a color-coded schematic according to a revised resource allocator routing table's matching normalized scores for a first processing node in the data center.

(17) FIG. 7B shows a table representing a revised resource allocator routing table of normalized scores for a second application in a first processing node in the data center.

(18) FIG. 7C shows a color-coded schematic according to a revised resource allocator routing table's matching normalized scores for a first processing node in the data center.

(19) FIG. 8A shows a first example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors in an intra-node implementation.

(20) FIG. 8B shows a second example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors for an inter-node implementation.

(21) FIG. 8C shows a third example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors in an intra-node implementation.

(22) FIG. 8D shows a fourth example schematic for calculating rated link latencies based on three hops between two reconfigurable processors in the plurality of reconfigurable processors in an intra-node implementation.

(23) FIG. 9 is a message sequence chart illustrating one implementation of compiling an application and generating an execution file for reconfigurable processors based on link bandwidth and link latency.

(24) FIG. 10 is a message sequence chart illustrating one implementation of a data processing system executing an application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(25) FIG. 11 is a message sequence chart illustrating one implementation of a data processing system executing a first application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(26) FIG. 12 is a message sequence chart illustrating one implementation of a data processing

system executing a second application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(27) FIG. 13 is a message sequence chart illustrating one implementation of a data processing system executing an application using a routing table for the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(28) FIG. 14 is a message sequence chart illustrating one implementation of a data processing system executing first and second applications using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(29) FIG. 15 is a message sequence chart illustrating one implementation of a data processing system executing first and second applications, such that first and second virtual routing tables are used for the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency.

(30) FIG. 16 is a system diagram illustrating a system including a host, a memory, and an example reconfigurable data processor on which the technology disclosed can be applied.

(31) FIG. 17 is a simplified block diagram of a top-level network and components of a CGRA (Coarse-Grained Reconfigurable Architecture).

(32) FIG. 18A is a simplified diagram of a tile and an array level network usable in the configuration of FIG. 17, where the configurable units are nodes on the array level network and are configurable to implement a Look-Up Table with input offsetting.

(33) FIG. 18B illustrates an example Switch unit connecting elements in an array level network.

(34) FIG. 19 is a block diagram illustrating an example configurable unit 1900, such as a Pattern Compute Unit (PCU), which is configured with the virtualization logic 1697.

(35) FIG. 20 is a block diagram illustrating an example configurable unit, such as a Pattern Memory Unit (PMU).

SUMMARY

(36) A data processing system is described that comprises a pool of reconfigurable dataflow resources and runtime logic. Reconfigurable dataflow resources in the pool of reconfigurable dataflow resources include a plurality of reconfigurable processors operatively coupled by links. The links have rated link bandwidths and rated link latencies.

(37) The runtime logic is operatively coupled to the pool of reconfigurable dataflow resources, and is configured to receive for a first application: (1) a first plurality of configuration files that contain configuration data; (2) a first configuration of a first plurality of virtual reconfigurable processors required to execute the first application, and virtual links between virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; and (3) a first specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors. The runtime logic allocates reconfigurable processors in the plurality of reconfigurable processors to the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors, and allocates links between the reconfigurable processors to the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors based on a link bandwidth comparison and a link latency comparison. The link bandwidth comparison compares the target link bandwidths specified by the first specification against the rated link bandwidths. The link latency comparison compares the target link latencies specified by the first specification against the rated link latencies. The runtime logic configures the allocated reconfigurable processors and the allocated links with configuration data in the first plurality of configuration files, and executes the first application using the configured reconfigurable processors and the configured links.

(38) In one implementation, the first configuration is encoded in a first virtual routing table that specifies a set of virtual links for each pair of virtual reconfigurable processors in the first plurality of virtual reconfigurable processors. In one implementation, the first specification is encoded in the first virtual routing table that specifies a target link bandwidth and a target link latency of each

virtual link in the set of virtual links.

(39) In one implementation, the runtime logic is further configured to receive for the first application a routing table that specifies a set of links for each pair of reconfigurable processors in the plurality of reconfigurable processors available to execute the first application.

(40) In one implementation, the routing table specifies a rated link bandwidth and a rated link latency of each link in the set of links.

(41) In one implementation, the runtime logic is further configured to allocate pairs of reconfigurable processors in the plurality of reconfigurable processors to pairs of virtual reconfigurable processors in the first plurality of virtual reconfigurable processors, and further configured to allocate links in the set of links to virtual links in the set of virtual links based on a link bandwidth comparison and a link latency comparison. The runtime logic compares target link bandwidths of the virtual links in the set of virtual links against rated link bandwidths of the links in the set of links. A link latency comparison by the runtime logic compares target link latencies of the virtual links in the set of virtual links against rated link latencies of the links in the set of links. The runtime logic configures the allocated pairs of reconfigurable processors and the allocated links with the configuration data in the first plurality of configuration files. The runtime logic executes the first application using the configured pairs of reconfigurable processors and the configured links.

(42) In one implementation, the runtime logic is further configured to receive for a second application: (1) a second plurality of configuration files that contain configuration data; (2) a second configuration of a second plurality of virtual reconfigurable processors required to execute the second application, and virtual links between virtual reconfigurable processors in the second plurality of virtual reconfigurable processors; and (3) a second specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors.

(43) In one implementation, the runtime logic is further configured to receive for the second application revised link bandwidths and revised link latencies of the allocated links. The revised link bandwidths and the revised link latencies result from allocation of parts of the rated link bandwidths and the rated link latencies to the first application. The runtime logic further allocates the allocated reconfigurable processors to the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors, and the allocated links to the virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors. The allocations are based on a link bandwidth comparison that compares the target link bandwidths specified by the second specification against the revised link bandwidths, and a link latency comparison that compares the target link latencies specified by the second specification against the revised link latencies. The runtime logic further configures the allocated reconfigurable processors and the allocated links with configuration data in the second plurality of configuration files, and executes the second application using the configured reconfigurable processors and the configured links.

(44) In one implementation, the first and second pluralities of configuration files, the first and second configurations, and the first and second specifications are generated by a compiler.

(45) In one implementation, the routing table is generated by a resource manager.

(46) In one implementation, the rated link latencies are rated based on a number of hops between two reconfigurable processors in the plurality of reconfigurable processors. In one implementation, a hop is an intermediate reconfigurable processor between the two reconfigurable processors. In one implementation, the intermediate reconfigurable processor is operative as a relay. In one implementation, the intermediate reconfigurable processor relay is operative bidirectionally. In one implementation, the intermediate reconfigurable processor relay is operative unidirectionally. In one implementation, a latency value is a time delay for traversing the intermediate reconfigurable processor relay. In one implementation, the latency value for traversing the intermediate

reconfigurable processor relay is an a priori known static value for calculating a normalized latency score. In one implementation, the latency delay value for traversing the intermediate reconfigurable processor relay is dependent on configuration routing and is a dynamic value for calculating the normalized latency scores.

(47) In one implementation, rated link latencies of a set of links between the two reconfigurable processors are normalized scores that are normalized across links in the set of links and thereby sum to a ceiling score. In one implementation, target link latencies of a set of virtual links between two virtual reconfigurable processors are normalized scores that are normalized across virtual links in the set of virtual links and thereby sum to the ceiling score.

(48) In one implementation, the rated link bandwidths are rated based on a hardware configuration of the links. In one implementation, the hardware configuration is a data transfer rate of the links. In one implementation, the hardware configuration is a link width of the links. In one implementation, the first and second applications are executed in parallel using the configured reconfigurable processors and the configured links.

(49) In one implementation, execution of the first application using the configured reconfigurable processors and the configured links further includes processing and streaming application data for the first application and outputs generated therefor using the configured reconfigurable processors and the configured links.

(50) In one implementation, execution of the second application using the configured reconfigurable processors and the configured links further includes processing and streaming application data for the second application and outputs generated therefor using the configured reconfigurable processors and the configured links.

(51) In one implementation, the allocated reconfigurable processors are on a same processing node. In one implementation, the allocated reconfigurable processors are on different processing nodes.

(52) In one implementation, the link bandwidth comparison precedes the link latency comparison. In one implementation, the link latency comparison precedes the link bandwidth comparison. In one implementation, the same unit is used to measure and compare the rated link bandwidths and the target link bandwidths. In one implementation, the same unit is used to measure and compare the rated link latencies and the target link latencies.

(53) A data processing system comprises a pool of reconfigurable dataflow resources and runtime logic. The reconfigurable dataflow resources in the pool of reconfigurable dataflow resources include a plurality of reconfigurable processors operatively coupled by links. The links have rated link bandwidths and rated link latencies.

(54) The runtime logic is operatively coupled to the pool of reconfigurable dataflow resources, and is configured to receive: (1) configuration files for first and second applications; (2) a first virtual routing table that specifies a first plurality of virtual reconfigurable processors required to execute the first application, and target link bandwidths and target link latencies of virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; and (3) a second virtual routing table that specifies a second plurality of virtual reconfigurable processors required to execute the second application, and target link bandwidths and target link latencies of virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors.

(55) The runtime logic allocates a set of reconfigurable processors in the plurality of reconfigurable processors to the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors and the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors, and links between reconfigurable processors in the set of reconfigurable processors to the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors and the virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors, based on a link bandwidth comparison and a link latency comparison. The link bandwidth comparison compares the rated link bandwidths

against the target link bandwidths specified by the first virtual routing table and the target link bandwidths specified by the second virtual routing table. The link latency comparison compares the rated link latencies against the target link latencies specified by the first virtual routing table and the target link latencies specified by the second virtual routing table. The runtime logic endeavors to assign reconfigurable processors and links between the reconfigurable processors assigned to each plurality of virtual reconfigurable processors such that target link bandwidths and target link latencies for each plurality of virtual reconfigurable processors meet their respective constraints in view of the rated link latencies and bandwidths.

(56) The runtime logic configures the allocated reconfigurable processors and allocated links with configuration data in the configuration files, and executes the first and second applications using the configured reconfigurable processors and the configured links.

(57) A data processing system comprises a pool of reconfigurable dataflow resources and runtime logic. The reconfigurable dataflow resources in the pool of reconfigurable dataflow resources include a plurality of reconfigurable processors operatively coupled by links. The links have rated link bandwidths and rated link latencies.

(58) The runtime logic is operatively coupled to the pool of reconfigurable dataflow resources, and is configured to receive: (1) the configuration files that define an application; (2) a configuration of virtual reconfigurable processors and virtual links between the virtual reconfigurable processors required to execute the application; and (3) a specification of target link bandwidths and target link latencies of the virtual links.

(59) The runtime logic allocates reconfigurable processors in the plurality of reconfigurable processors to the virtual reconfigurable processors, and allocates links between the reconfigurable processors to the virtual links. The allocations are based on a link bandwidth comparison that compares the target link bandwidths against the rated link bandwidths, and a link latency comparison that compares the target link latencies against the rated link latencies.

(60) The runtime logic configures the allocated reconfigurable processors and the allocated links with configuration data in the configuration files, and executes the application using the configured reconfigurable processors and the configured links.

(61) A system is described that has a node and runtime logic. The node has a plurality of processing elements operatively coupled by interconnects. The runtime logic is configured to receive target interconnect bandwidth, target interconnect latency, rated interconnect bandwidth and rated interconnect latency. The runtime logic responds by allocating to configuration files defined by the application graph: (1) processing elements in the plurality of processing elements, and (2) interconnects between the processing elements. The runtime logic further responds by executing the configuration files using the allocated processing elements and the allocated interconnects.

(62) In one implementation, the processing elements are respective arrays of configurable units. In one implementation, the processing elements are respective tiles configured on an array of configurable units, where each of the tiles comprises a non-overlapping group of configurable units in the array of configurable units.

(63) In one implementation, the processing elements are respective tiles configured across arrays of configurable units, where each of the respective tiles comprises a non-overlapping group of configurable units in a corresponding array of configurable units in the arrays of configurable units. In one implementation, the processing elements are respective configurable units of an array of configurable units. In one implementation, configurable units in the respective configurable units are Pattern Compute Units (PCUs) and Pattern Memory Units (PMUs).

(64) A computer-implemented method is described that includes: (1) allocating processing elements and interconnects between the processing elements to applications requesting execution based on achievable interconnect bandwidth and achievable interconnect latency of the interconnects to satisfy target interconnect bandwidth and target interconnect latency requested by the applications; and (2) executing the applications using the allocated processing elements and the allocated

interconnects.

(65) A computer-implemented method is described that includes: (1) allocating processing elements and interconnects between the processing elements to applications requesting execution based on rated bandwidth and rated latency of the interconnects to satisfy target bandwidth and target latency of the applications; and (2) executing the applications using the allocated processing elements and the allocated interconnects.

(66) A computer-implemented method is described that includes: (1) allocating processing elements to applications requesting execution based on data transfer rate between the processing elements; and (2) executing the applications using the allocated processing elements.

(67) In one implementation, the node topology of reconfigurable processors has a Dragonfly cluster implementation providing point-to-point linking in an all-to-all connections using direct links (i.e., without an external switching fabric).

(68) In one implementation, the node topology of reconfigurable processors has a 3D Torus cluster implementation providing point-to-point linking in an all-to-all connections using direct links.

(69) In one implementation, the node topology of reconfigurable processors has a Fat Tree cluster implementation providing point-to-point linking connections using reconfigurable processors links and intermediate reconfigurable processors links. In one implementation, the node topology of reconfigurable processors has a Hypercube cluster implementation providing point-to-point linking in an all-to-all connections using direct links.

(70) In one implementation, the node topology of reconfigurable processors has a HyperX cluster implementation providing point-to-point linking in an all-to-all connections using direct links.

(71) One or more implementations of the technology disclosed, or elements thereof can be implemented in the form of a computer product including a non-transitory computer readable storage medium with computer usable program code for performing the method steps indicated. Furthermore, one or more implementations of the technology disclosed, or elements thereof can be implemented in the form of an apparatus including a memory and at least one processor that is coupled to the memory and operative to perform exemplary method steps. Yet further, in another aspect, one or more implementations of the technology disclosed or elements thereof can be implemented in the form of means for carrying out one or more of the method steps described herein; the means can include (i) hardware module(s), (ii) software module(s) executing on one or more hardware processors, or (iii) a combination of hardware and software modules; any of (i)-(iii) implement the specific techniques set forth herein, and the software modules are stored in a computer readable storage medium (or multiple such media).

(72) These and other features, aspects, and advantages of the technology disclosed will become apparent from the following detailed description of illustrative implementations thereof, which is to be read in connection with the accompanying drawings.

DETAILED DESCRIPTION

(73) The following discussion is presented to enable any person skilled in the art to make and use the technology disclosed and is provided in the context of a particular application and its requirements. Various modifications to the disclosed implementations will be readily apparent to those skilled in the art, and the general principles defined herein may be applied to other implementations and applications without departing from the spirit and scope of the technology disclosed. Thus, the technology disclosed is not intended to be limited to the implementations shown but is to be accorded the widest scope consistent with the principles and features disclosed herein.

(74) Data Center

(75) Systems and processes for providing Reconfigurable Processor-as-a-Service (RPaaS) are described. The systems and processes will be described with reference to FIG. 1 showing an architectural level schematic of a node within a data center **100** in accordance with an implementation. Because FIG. 1 is an architectural diagram, certain details of the data center **100**

are intentionally omitted to improve the clarity of the description. It may be noted that data center **100** can include the same, more, or fewer elements configured in the same or different manner in other implementations. The discussion of FIG. **1** will be organized as follows. First, the elements of the figure will be described, followed by their interconnections. Then, the use of the elements in the system will be described in greater detail.

(76) FIG. **1** shows elements of a first processing node of a data center **100**. In FIG. **1**, the processing node is configured to collaboratively execute configuration files for applications in a distributed fashion. One skilled in the art will appreciate that the data center **100** can have any number of processing nodes operatively coupled for data communications through a network **136** (also called herein “network fabric **136**”). Examples of the network **136** include a Storage Area Network (SAN) and a Local Area Network (LAN). The SAN can be implemented with a variety of data communications fabrics, devices, and protocols. For example, the fabrics for the SAN can include Fibre Channel, Ethernet, InfiniBand, Serial Attached Small Computer System Interface (‘SAS’), or the like. Data communications protocols for use with the SAN can include Advanced Technology Attachment (‘ATA’), Fibre Channel Protocol, Small Computer System Interface (‘SCSI’), Internet Small Computer System Interface (‘iSCSI’), HyperSCSI, Non-Volatile Memory Express (‘NVMe’) over fabrics, or the like.

(77) The LAN can also be implemented with a variety of fabrics, devices, and protocols. For example, the fabrics for the LAN can include Ethernet (**802.3**), wireless (**802.11**), or the like. Data communication protocols for use in the LAN can include Transmission Control Protocol (‘TCP’), User Datagram Protocol (‘UDP’), Internet Protocol (IP), Hypertext Transfer Protocol (‘HTTP’), Wireless Access Protocol (‘WAP’), Handheld Device Transport Protocol (‘HOTP’), Session Initiation Protocol (‘SIP’), Real-time Transport Protocol (‘RTP’), or the like.

(78) The network **136** also connects other network components in the data center **100**. Examples of other network components include links, switches, routers, load balancers, hypervisors, and Application Programming Interfaces (APIs). Along the network **136**, the switches, for example, can receive packets via a plurality of input ports and can transmit packets via a plurality of output ports. The processing nodes in the data center **100** can communicate with each other through the network **136** using a variety of networking paths established by the switches. Another example of the network **136** is a Wide Area Network (WAN).

(79) A processing node (or node) is an addressable application running on a hardware device or virtual device that attaches to a network, and is capable of sending, receiving, or forwarding information over a communications channel to or from other processing nodes. Examples of electronic devices which can be deployed as hardware processing nodes include all varieties of computers, workstations, laptop computers, handheld computers, and smartphones. Processing nodes can be implemented in a cloud-based server system. More than one virtual device configured as a processing node can be implemented using a single physical device.

(80) The data center **100** comprises a pool of reconfigurable dataflow resources. The pool of reconfigurable dataflow resources can have a variety of compute scales and hierarchies. The pool of reconfigurable dataflow resources can be a single processing node operatively coupled to a plurality of reconfigurable processors, which in turn is supported by different bus and memory resources. The processing node can have a host processor (e.g., a CPU) that exchanges data with the reconfigurable processors, for example, over a local bus like Peripheral Component Interconnect Express (PCIe) interface. The host processor can have a runtime processor (or a runtime logic) that manages resource allocation, memory corresponding, and execution of configuration files for applications requesting execution from the host processor.

(81) The pool of reconfigurable dataflow resources can be a rack (or cluster) of processing nodes connected through the network **136**. Each processing node in the rack can run a respective plurality of reconfigurable processors and include a respective host processor configured with a respective runtime processor. The runtime processors, distributed across the processing nodes, communicate

with each other to provide unified access to reconfigurable processors attached not only to their own processing node but also to reconfigurable processors attached to every other processing node of the data center **100**.

(82) The pool of reconfigurable dataflow resources can be a pod that comprises a plurality of racks connected through the network **136**. The pool of reconfigurable dataflow resources can be a superpod that comprises a plurality of pods connected through the network **136**. The pool of reconfigurable dataflow resources can be a zone that comprises a plurality of superpods connected through the network **136**. The pool of reconfigurable dataflow resources can be the data center **100** that comprises a plurality of zones connected through the network **136**.

(83) The pool of reconfigurable dataflow resources can include link (interconnect or transfer) and bus resources. Examples of the link resources include PCIe channels. Examples of bus resources include Direct Memory Access (DMA) channels, and Double Data Rate (DDR) channels. The pool of reconfigurable dataflow resources can include memory (or storage) resources. Examples of the memory resources include main memory (e.g., off-chip/external Random-Access Memory (RAM), Dynamic Random-Access Memory (DRAM), Static Random-Access Memory (SRAM), NAND flash), local secondary storage (e.g., local disks (e.g., Hard-Disk Drive (HDD), Solid-State Drive (SSD))), and remote secondary storage (e.g., distributed file systems, web servers). Other examples of the memory resources include latches, registers, flops, bypass networks, and caches (e.g., ones explicitly addressed by RAMs/DRAMs/SRAMs). The pool of reconfigurable dataflow resources is dynamically scalable to meet the performance requirements of applications requesting execution. The applications access the pool of reconfigurable dataflow resources over one or more networks (e.g., the Internet).

(84) The discussion now returns to the first processing node of the data center **100**. The first processing node comprises a host processor **102**. Examples of the host processor **102** include x86 and x64 processors. The first host processor **102** interfaces with a host memory **150** (e.g., RAM). The first host processor **102** has a compiler **112** to compile applications and runtime logic **140** to execute the compiled applications on a plurality of reconfigurable processors (e.g., RPs **142**). The runtime logic **140** is configured to provide on-demand access to the pool of reconfigurable dataflow resources, which can be rapidly provisioned and released with minimal management effort or service provider interaction.

(85) Examples of the reconfigurable processors (e.g., RPs **142**) include Field Programmable Gate Arrays (FPGAs), Coarse-Grained Reconfigurable Architectures (CGRAs), Application-Specific Integrated Circuits (ASICs), and Application Specific Instruction-set Processor (ASIP). The reconfigurable processors (e.g., RPs **142**) interface with a reconfigurable processor memory **162** (e.g., DRAM). Each of the reconfigurable processors (e.g., RPs **142**) includes an array of configurable units (e.g., compute units and memory units) in a programmable interconnect fabric. The array of configurable units in a reconfigurable processor is partitionable into a plurality of subarrays (or tiles) of configurable units. Additional details about one implementation of the architecture of the reconfigurable processors are discussed later in this application. In other implementations, the processing nodes in the data center **100** include processors instead of/in addition to the reconfigurable processors (e.g., RPs **142**). Examples of such processors include Graphics Processing Units (GPUs) and Digital Signal Processors (DSPs). In other implementations, the first processing node does not comprise a host processor **102**.

(86) A Network Interface Controller **132** (e.g., NIC, SmartNIC) connects the first host processor **102** and the reconfigurable processors (e.g., RPs **142**) to the Network Fabric **136** via a switching fabric **110**. The Network Interface Controller **132** interfaces to the Network Fabric **136** via link **130**. Examples of the network fabric links are Peripheral Component Interconnect Express (PCIe), Cache Coherent Interconnect for Accelerators (CCIX), Compute Express Link (CXL), Fibre Channel, Ethernet, InfiniBand, and Open Coherent Accelerator Processor Interface (OpenCAPI).

(87) A Switching Fabric **110** (e.g. PEX) interfaces to the first host processor **102**, the reconfigurable

processors (e.g., RPs **142**) and to the Network Interface Controller **132**. The example Switching Fabric **110** uses eight point-to-point local links **120-127** operatively coupling to the reconfigurable processors (e.g., RPs **142**), a ninth point-to-point link **128** operatively coupling to a first host processor **102** and a tenth point-to-point local link **129** to couple to the Network Interface Controller **132**. In some implementations, the Switching Fabric **110** supports 16, 24, 32 or more links. The port count of a Switching Fabric **110** refers to the number of links it can support. So, an eight port Switching Fabric supports eight links. Examples of links are Peripheral Component Interconnect Express (PCIe), Cache Coherent Interconnect for Accelerators (CCIX), Compute Express Link (CXL), Fibre Channel, Ethernet, InfiniBand, and Open Coherent Accelerator Processor Interface (OpenCAPI). The Network Interface Controller **132** interfaces to the Network Fabric **136** via link **130**.

(88) The reconfigurable processors (e.g., RPs **142**) topology comprises links providing an all-to-all linking via the Switching Fabric **110**. Additional point-to-point links shown in FIG. **1** between reconfigurable processors (e.g., RPs **142**) do not require the bridging of a switching fabric **110**. These links include cross coupling and direct linking between four reconfigurable processors. In some implementations, direct linking of reconfigurable processors does not require the bridging of a switching fabric **110**. In some implementations, direct linking of reconfigurable processors is accomplished for five or more reconfigurable processors (e.g., RPs **142**). In some implementations, a Dragonfly cluster technology implementation can provide point-to-point linking in an all-to-all topology for reconfigurable processors. In one implementation, a Fat Tree cluster implementation can be used. In other implementations, a 3D Torus cluster implementation can be used. In other implementations, a Hypercube or HyperX cluster technology can be implemented. One skilled in the art will appreciate that direct linking reconfigurable processors (e.g., RPs **142**) is dependent on the reconfigurable processor pin count, printed circuit board (PCB) routing considerations, as well as cooling and power consumption to support a node in an all-to-all topology with direct links.

(89) Having described the elements and interconnections of FIG. **1**, the discussion now turns to determining link bandwidth and link latency for efficient streaming of data between processors on a same processing node and on different processing nodes and the sharing of links thereof.

(90) Link Bandwidth and Link Latency

(91) FIGS. **2A**, **2B** and **2C** **200** compare target link bandwidths to rated link bandwidths of a first processing node in the data center of FIG. **1** using tables and a color-coded flattened representation of a plurality of reconfigurable processors (e.g., RPs **142**). The processing elements RP0 **242**, RP1 **243**, RP2 **244**, and RP3 **245** are reconfigurable processors in a first processing node.

(92) FIG. **2A** shows a table representing an example specification for target link bandwidths of a first processing node in the data center of FIG. **1**. The table specifies two rows of virtual links listing target link bandwidths for pairs of virtual reconfigurable processors. The specification lists bandwidths for the three virtual pairs. The first virtual pair of reconfigurable processors corresponds to the pair RP0 **242** to RP1 **243**. The second virtual pair corresponds to the pair RP0 **242** to RP2 **244** and the third virtual pair corresponds to the pair RP0 **242** to RP3 **245**.

(93) FIG. **2B** shows a table representing an example specification for rated link bandwidths of a first processing node in the data center of FIG. **1**. The rated link bandwidths table has three rows of rated links available for the three virtual pairs. The three rows of rated links do not provide bandwidth to every virtual pair. Link **2** is not available for the second virtual pair corresponding to RP0 **242** to RP2 **244**. Link **3** is not available for the two pairs RP0 **242** to RP1 **243** and RP0 **242** to RP2 **244**. Only RP0 **242** to RP3 **245** has three rated links available.

(94) FIG. **2C** shows an example color-coded flattened schematic comparing target link bandwidths to rated link bandwidths for the processing elements RP0 **242**, RP1 **243**, RP2 **244**, and RP3 **245** which are four reconfigurable processors (e.g., RPs **142**) of a first processing node in the data center of FIG. **1**. A graphical comparison of target link bandwidths to rated link bandwidths has color coding. The rated link bandwidth is in black font and the target link bandwidth is in orange

font. For example, the RP0-RP1 Link2 with a rated link bandwidth of 5 GB/s in the black font is less than the virtual specification of 15 GB/s in orange. Therefore, since the RP0-RP1 Link2 does not meet the virtual specification, the line drawn between RP0 242 to RP1 243 is in the color red. The first virtual pair corresponds to the pair RP0 242 to RP1 243. The second virtual pair corresponds to the pair RP0 242 to RP2 244, and the third virtual pair corresponds to the pair RP0 242 to RP3 245. For the first virtual pair, only the RP0-RP1 Link1 whose line is in blue satisfies the virtual request of the first processing node. For the second virtual pair, only the PR0-PR2 Link1 meets the requirements of Virtual Link2 and is color-coded blue. The third virtual pair meets all the virtual specifications. The best link is in the green color for RP0 242 to RP3 245, where the rated Link3 corresponds to the Virtual Link1 and the rated Link1 in the color blue corresponds to Virtual Link2.

(95) A bandwidth equation summarizes the requirements to meet when running applications concurrently over shared links between RPs i and j:

(96) $BW_k(ij) > BW_k(a)(ij) + BW_k(b)(ij)$

Where: $BW_k(ij)$ —Total bandwidth between nodes i and j through path k $BW_k(a)(ij)$ —Bandwidth requirement from a first application (a) between nodes i and j $BW_k(b)(ij)$ —Bandwidth requirement from a second application (b) between nodes i and j

If a third application $BW_k(c)(ij)$ is presented, the Kernel Resource Manager (KRM) checks each set of links k between the two nodes (i, j) that satisfy:

(97) $BW_k(ij) > [BW_k(a)(ij) + BW_k(b)(ij) + BW_k(c)(ij)]$ with the most optimum routing.

(98) In some implementations, the rated link bandwidths are rated based on a hardware configuration of the links. In some implementations, the hardware configuration is a data transfer rate of the link. In some implementations, the hardware configuration is a link width of the links. An example of PCIe v5.0 link bandwidth is 32 GB/s for each lane/way. A 4 lane/way PCIe v5.0 implementation would supply a total bandwidth of 128 GB/s. An example of PCIe v6.0 link bandwidth is 64 GB/s for each lane/way. A 4 lane/way PCIe v6.0 implementation would supply a total bandwidth of 256 GB/s. An example InfiniBand link bandwidth for a 12-way HDR implementation is 600 GB/s. The InfiniBand Trade Association (IBTA) roadmap shows future InfiniBand link bandwidths increasing to NDR having 1.2 TB/s, XDR having 2.4 TB/s, and GDR having 4.8 TB/s.

(99) Having described the aspect of comparing bandwidths, the discussion now turns to choosing link latency for efficiently streaming of data concurrently for a first and for a second application over shared links by allocating link latency for the first and second applications between processors on an intra or inter processing node.

(100) FIG. 3 shows an example compiler assignment 300 of rated link latencies in a color-coded schematic as normalized scores from one to one hundred for a first processing node in the data center of FIG. 1. The runtime assignment of Rated Link (RL) latency is a normalized score from one to one hundred and assigned to every link. For example, the link RP0 242 to RP1 243 is corresponding to RL1[01] with a normalized score of 20. This figure illustrates all combinations to help identify the shortest path from the perspective of RP0 242. Calculating rated link latency is not only from the first reconfigurable processor (e.g., RP0), but calculates from all the reconfigurable processors in each node creating many permutation tables from RP1's perspective, RP2's perspective, RP3's perspective, etc. The Kernel Resource Manager (KRM) manages links in order to meet latency and bandwidth specifications when allocating reconfigurable processors to reconfigurable processor pairs from the data resource pool. When an application specifies virtual resources for matching computational and memory requirements, the KRM identifies links connecting and allocating data resources by comparing the latency costs.

(101) FIGS. 4A, 4B and 4C 400 depict in table format and graphically target link latency and routing table links as normalized scores.

(102) FIG. 4A shows an example table representing a first specification for target link latency as

normalized scores of a first processing node in the data center of FIG. 1. The table specifies two rows of virtual links listing target link latency for virtual pairs of reconfigurable processors. The specification lists normalized scores for the three virtual pairs. The first virtual pair corresponds to the pair **RP0 242** to **RP1 243**. The second virtual pair corresponds to the pair **RP0 242** to **RP2 244** and the third virtual pair corresponds to the pair **RP0 242** to **RP3 245**.

(103) FIG. 4B shows an example table representing a resource allocator routing table numbered as normalized scores for a first processing node in the data center of FIG. 1. The table specifies two rows of (physical) links listing rated latency for pairs of reconfigurable processors. The routing table lists normalized scores for the three pairs where the first pair corresponds to the pair **RP0 242** to **RP1 243**, the second pair is the pair **RP0 242** to **RP2 244** and the third pair is the pair **RP0 242** to **RP3 245**.

(104) FIG. 4C shows an example color-coded schematic comparing link bandwidths and link latencies for matching the normalized scores of target link latency to resource allocator routing table's normalized scores of rated link latency in a first processing node in the data center of FIG. 1. The target latency is in a blue font and the routing table normalized value is in green font. For example, the link between **RP0 242** and **RP1 243**, named **RP0-RP1 Link1**, shows the target link latency 1 value of 20 taken from FIG. 4A is in blue. The rated link 1 with a normalized value of 60 in green is taken from FIG. 4B. Likewise, the link between **RP0 242** and **RP1 243** named **RP0-RP1 Link2** shows a target link latency of 80 in blue from the table in FIG. 4A, and a rated link latency of 40 in green from the routing table in FIG. 4B. Since these values do not route in the most optimum way, matching the normalized score with the least latency to the link with the highest latency target, the lines are both shown in red color. The links shown in blue represent an optimal routing where the normalized scores match relatively closely. For example, the default routing between **RP0 242** and **RP3 245** is the link **RP0-RP3 Link2** with a value 40 in blue for target link latency and a value of 30 in green from the routing table of FIG. 4B. This example shows the default matching of the lower rated latency in the routing table to the lower target link latency. The next figure shows allocation corrections to default values between **RP0 242** and **RP1 243**.

(105) FIGS. 4A, 5A and 5B 500 depict in table format and graphically target link latency and a revised routing table of normalized scores.

(106) FIG. 4A repeats in 500 and shows a table representing a first specification for target link latency as normalized scores for a first processing node in the data center of FIG. 1. The table specifies two rows of virtual links listing target link latency for virtual pairs of reconfigurable processors. The specification lists normalized scores for the three virtual pairs. The first virtual pair corresponds to the pair **RP0 242** to **RP1 243**. The second virtual pair corresponds to the pair **RP0 242** to **RP2 244** and the third virtual pair corresponds to the pair **RP0 242** to **RP3 245**.

(107) FIG. 5B shows a table representing a revised resource allocator routing table of normalized scores for a first processing node in the data center of FIG. 1. Continuing with the same color coding as shown in FIG. 4C, the normalized values of target link latency are in a blue font and the routing table's normalized values are in a green font. To correct the routing problem disclosed in FIG. 4C, of two links between **RP0 242** and **RP1 243**, **RP0-RP1 Link1** swaps with **RP0-RP1 Link2**. Referring to the original target link latency of FIG. 4A where the target link latency 1 is a normalized value of 20 and the target link latency 2 has a value 80. Now, after swapping, FIG. 5A shows the target link latency 20 is matched to **RP0-RP1 Link2** which has a rated link latency 40, while the target link latency 80 is allocated with **RP0-RP1 Link1** having a rated link latency 60.

(108) FIG. 5A shows a color-coded schematic according to a revised resource allocator routing table of normalized scores for a first processing node in the data center of FIG. 1. All links are now in blue after allocation corrections to the two **RP0-RP1** links.

(109) FIGS. 6A, 6B and 6C 600 depict in table format and graphically a second set of target link latencies and routing table links as normalized scores.

(110) FIG. 6A shows a table representing a second specification for target link latency as

normalized scores for a first processing node in the data center of FIG. 1. The table specifies two rows listing target link latency for virtual pairs of reconfigurable processors. There are three columns listing the virtual pairs, a first column VRP0 to VRP1, a second column VRP0 to VRP2, and a third VRP0 to VRP3. The specification lists normalized scores for the three virtual pairs. The first virtual pair corresponds to the pair RP0 242 to RP1 243. The second virtual pair corresponds to the pair RP0 242 to RP2 244 and the third virtual pair corresponds to the pair RP0 242 to RP3 245.

(111) FIG. 6B, which is FIG. 4B repeated in 600 and renamed FIG. 6B, shows a table representing a resource allocator routing table of normalized scores of a first processing node in the data center of FIG. 1. The table specifies two rows of (physical) links listing rated latency scores: rated link1 score and rated link2 score. The routing table lists normalized scores for the three pairs of reconfigurable processors. The first pair corresponds to the pair RP0 242 to RP1 243. The second pair corresponds to the pair RP0 242 to RP2 244 and the third pair corresponds to the pair RP0 242 to RP3 245.

(112) FIG. 6C shows an example color-coded schematic comparing virtual pairs for matching the normalized scores of target link latency to resource allocator routing table normalized scores of rated latency in a first processing node in the data center of FIG. 1. The target latency is in a blue font and the routing table normalized value is in green. For example, the link between RP0 242 and RP1 243 is named RP0-RP1 Link1 and shows the target link latency 1 value of 30 in blue from FIG. 6A. The rated link latency 1 with a value of 60 in FIG. 6B is shown in green. Likewise, another link between RP0 242 and RP1 243 is named RP0-RP1 Link2 and shows a target link latency of 70 in blue taken from the specification in FIG. 6A. The RP0-RP1 Link2 shows a rated link latency of 40 in green taken from the routing table in FIG. 6B. Since these values do not route in the most optimum way, matching the link with least rated latency of normalized score to the virtual link with the highest latency target, the lines are both shown in RED. The links shown in blue represent an optimal routing where the normalized scores match relatively closely. For example, the default routing between RP0 242 and RP3 245 is the link RP0-RP3 Link2 with a rated latency value 30 in green from the routing table of FIG. 6B allocated for the target link latency value 45 in blue. This example shows the default matching of the lower target link latency to the lower link latency in the routing table.

(113) One skilled in the art will appreciate that values in the specification of FIG. 6A and values in the table of FIG. 6B are copied into FIG. 6C using the color codes of blue and green. The next FIG. 700 shows allocations and corrections to default values of the link between RP0 242 and RP1 243.

(114) FIG. 7B shows a table representing a revised resource allocator routing table of normalized scores for a first processing node in the data center of FIG. 1. Continuing with the same color coding as shown in FIG. 4C, the target latency normalized values are in a blue font and the routing table normalized values of FIG. 7B are in a green font. To correct the routing problem disclosed in FIG. 6C, the link RP0 242 to RP1 243 named RP0-RP1 Link1 swaps with RP0-RP1 Link2 for the original target link latency 1 of FIG. 6A, where the original target link latency 1 is a normalized value of 30 and the original target link latency 2 is a value of 70.

(115) FIG. 7A shows a color-coded schematic according to a revised resource allocator routing table matching the normalized scores of target link latency to the normalized scores of the routing table for a first processing node in the data center of FIG. 1. Now, as seen in the FIG. 7A, the virtual link of target latency 30 is matched to RP0-RP1 Link2 with a rated value of 40, and target link latency of 70 is allocated with RP0-RP1 Link1 having a rated latency value of 60. All links are now in blue after allocation corrections to the two RP0-RP1 links.

(116) The concept of counting hops as a factor in latency calculations is disclosed in the following FIGS. 8A, 8B, 8C, and 8D. The normalized scores include calculations of the required hops and bandwidths available. Some switches incur 100 nanosecond delays to traverse. Tables listing delay values for various network elements are tabulated according to the hop counts and specific device

delays. The hop counts shown are only between reconfigurable processors. It does not include hop counts incurred by host processors.

(117) FIG. **8A** shows a first example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors for an intra-node implementation. A first hop is from the reconfigurable processor **RP0** to **RP7** via the link **121** through the switching fabric **110** to link **127**. A second hop is from **RP7** to **RP1** via link **127** through the switching fabric **110** to link **122**. The links are in green color.

(118) FIG. **8B** shows a second example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors for an inter-node implementation. A first hop is from the reconfigurable processor **RP0** of Node **1** connecting to **RP7** in Node **2** via the link **121** through the switching fabric **110** to link **127**. A second hop is the connection of **RP7** in Node **2** to **RP1** in Node **1** via link **127** through the switching fabric **110** to link **122**. In some implementations the switching fabric can support one node, two nodes, three nodes, four nodes or more nodes.

(119) FIG. **8C** shows a third example schematic for calculating rated link latencies based on two hops between two reconfigurable processors in the plurality of reconfigurable processors in an intra-node implementation, where a hop is an intermediate reconfigurable processor between the two reconfigurable processors. In some implementations, the reconfigurable processor **RP4** is operative as a relay in a unidirectional direction. In some implementations, the reconfigurable processor **RP4** is operative as a relay in a bi-directional direction. In some implementations, the reconfigurable processor **RP4** is operative with an a priori static delay for calculating the normalized latency scores. In some implementations, the reconfigurable processor **RP4** is operative with a dynamic delay based on configuration files and routing when calculated for the normalized latency scores.

(120) FIG. **8D** shows a fourth example schematic for calculating rated link latencies based on three hops between two reconfigurable processors in the plurality of reconfigurable processors in, for example, an intra-node implementation. A first hop is from the reconfigurable processor **RP0** to **RP7** via the link **121** through the switching fabric **110** to link **127**. A second hop is from **RP7** to **RP1** via link **127** through the switching fabric **110** to link **122**. A third hop is from the reconfigurable processor **RP1** to **RP5** via the link **122** through the switching fabric **110** to link **126**. The links are in green color.

(121) One skilled in the art will appreciate that the number of hops counted for the latency normalization score is architecture dependent based on the permutations of traversing links, traversing a plurality of switching fabrics, reconfigurable processor relays, switches, routers, and bandwidth available. In some implementations, some links can be faster because of internal routing of the reconfigurable processors. The Kernel Resource Manager searches out the fastest links according to a knapsack or greedy algorithm, meeting the application specifications' input to the compiler. In some implementations, an Open Shortest Path First (OSPF) is a routing protocol for Internet Protocol (IP) networks. It uses a Link State Routing (LSR) algorithm that falls into the group of Interior Gateway Protocols (IGPs), operating within a single Autonomous System (AS) and defined as OSPF Version 2 in RFC 2328 (1998) for IPv4.

(122) Having described allocating link latency for first and second applications running concurrently over the same links for efficient streaming of data between processors on a same processing node and on different processing nodes and the sharing of links thereof, the discussion now turns to the aspects of compiling an application and the generation of execution files disclosed by the technology disclosed.

(123) Compiling Applications and Generating Execution Files

(124) FIG. **9** is a message sequence chart **900** showing the loading of an application **910** into a compiler **112** and the generation of configuration files. The runtime logic **140** also comprises allocation logic, as explained supra. The compiler **112** compiles an application **910** to generate a

graph that includes one or more configuration files corresponding to the application **910**. The runtime logic **140** receives an application **910** and execute one or more configuration files by sending a Graph.PEF (PEF stands for “pattern execution file”) to one or more of the reconfigurable processors (e.g., RPs **142**). The reconfigurable processors (e.g., RPs **142**) process the configuration files and send the generated outputs to the first host processor **102** using at least one of the reconfigurable processors-to-host processors. In some implementations, the node does not require or comprise a host processor.

(125) At operation one, the application **910** sends an application file, also referred to as a graph, to a compiler **112**. A series of compiler operations will generate an execution file and a plurality of configuration files. A compiler **112** compiles application(s) **910** and generates configuration files. The configuration files specify virtual dataflow resources required to execute the application graph.

(126) At operation **2a**, the compiler **112** generates configuration files, including a configuration of virtual reconfigurable processors (RPs). At operation **2b**, the compiler generates a configuration of virtual links between RPs. At operation **2c**, the compiler generates a target link bandwidth specification. At operation **2d**, the compiler generates a target latency specification. At operation **2E**, the compiler generates and places the metadata of operations **2a-2d** into an execution file. At operation three, the compiler sends the execution file (Graph.PEF) to the runtime logic **140** for execution on a first processing node in the data center of FIG. **1**. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. **9**. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. **9**, operations three can comprise streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(127) Loading and Executing Applications

(128) FIG. **10** is a message sequence chart **1000** illustrating one implementation of a data processing system executing an application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency. A runtime logic **1010** compares target link latency and target link bandwidth specifications to rated link latency and rated link bandwidth values. Based on target specifications meeting rated values, the system allocates dataflow resources. The runtime logic **1010** allocates pairs of reconfigurable processors and corresponding links for generating the configuration files used to load and execute an application.

(129) At operation one, runtime logic **1010** receives an execution file comprising configuration files that define an application. At operation two, configuration files, including the configuration of virtual reconfigurable processors and virtual links defining the application, is processed by the runtime logic **1010**. At operation three, the runtime logic **1010** compares the rated link bandwidth to target link bandwidth of the virtual links defined by the specification. At operation four, the runtime logic **1010** compares the rated link latency to target link latency of the virtual links specification. At operation five, the runtime logic **1010** configures and allocates pairs of reconfigurable processors and links according to the specification. At operation six, the runtime logic **1010** loads reconfigurable processors (e.g., RPs **142**) with configuration files and executes the application using the configured pairs of reconfigurable processors and the configured links on a first processing node in the data center of FIG. **1**.

(130) In some implementations, the operations three and four are reversed with the comparison of link latency preceding the comparison of link bandwidth. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. **10**. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. **10**, operation six comprises streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control

Protocol (TCP). In some implementations, operation six can apply to application graphs which are executed in parallel and/or across different processing nodes.

(131) FIG. **11** is a message sequence chart **1100** illustrating one implementation of a data processing system executing a first application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency. A runtime logic **1110** compares a target link latency and target link bandwidth specification to rated link latency and rated link bandwidth values. Based on rated system dataflow resources available to meet target specifications, the runtime logic **1110** allocates pairs of reconfigurable processors and corresponding links for generating the configuration files used to load and execute an application.

(132) At operation one, an execution file comprising configuration files that define an application is received by the runtime logic **1110**. At operation two, the specification defining the application is processed by the runtime logic **1110**, where the execution file provides a first configuration of virtual reconfigurable processors and virtual links. At operation three, the runtime logic **1110** processes the first configuration encoded in a first virtual routing table that specifies a set of virtual links for each pair of virtual reconfigurable processors in the first plurality of virtual reconfigurable processors. At operation four, the first specification is encoded in the first virtual routing table that specifies a target link bandwidth for each virtual link in the set of virtual links. At operation five, the first specification is encoded in the first virtual routing table that specifies a target link latency of each virtual link in the set of virtual links. At operation six, the runtime logic **1110** is further configured to receive, for the first application, a routing table that specifies a set of links for each pair of reconfigurable processors in the plurality of reconfigurable processors available to execute the first application. At operation seven, the routing table specifies a rated link bandwidth of each link in the set of links. At operation eight, the routing table specifies a rated link latency of each link in the set of links. At operation nine, the runtime logic **1110** compares the rated link bandwidth to target link bandwidth of the virtual links defined by the specification. At operation ten, the runtime logic **1110** configures the allocated pairs of reconfigurable processors and the allocated links with the configuration data in the first configuration files. At operation eleven, the runtime logic **1110** loads and executes the first application using the configured pairs of reconfigurable processors and the configured links on a first processing node in the data center of FIG. **1**.

(133) In some implementations, the specification is included in metadata header files. In some implementations, the operations seven and eight are reversed with operation eight preceding operation seven. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. **11**. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. **11**, operation eleven comprise streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(134) FIG. **12** is a message sequence chart **1200** illustrating one implementation of a data processing system executing a second application using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency. A runtime logic **1210** compares a target link latency and target bandwidth specification to revised rated link latency and revised rated bandwidth values. Based on rated system dataflow resources available to meet target specifications, the runtime logic **1210** allocates pairs of reconfigurable processors and corresponding links for generating the configuration files used to load and execute an application.

(135) At operation one, a compiler sends a second execution file comprising configuration files that define an application. At operation two, the runtime logic **1210** receives the second set of configuration files. At operation three, the second configuration of virtual reconfigurable processors, virtual links, target link latencies and target link bandwidths defining the application is processed by the runtime logic **1210**. At operation four, the runtime logic **1210** receives a revised rated link bandwidth, and a revised link latency in a second target link bandwidth and target link

latency specification. At operation five, the runtime logic **1210** compares the revised rated link bandwidth to the second target link bandwidth defined by the second specification and compares the revised rated link latency to the second target link latency defined by the second specification. At operation six, the runtime logic **1210** configures and allocate pairs of reconfigurable processors and links according to the second specification. At operation seven, the runtime logic **1210** loads reconfigurable processors (e.g., RPs **142**) with configuration files and executes the application on a first processing node in the data center of FIG. **1**.

(136) Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. **12**. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. **12**, operation seven comprises streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(137) Having described executing and generating applications over the same links for efficient streaming of data between processors on a same processing node and on different processing nodes, the discussion now continues with the respect to creating a routing table for an application.

(138) Kernel Resource Manager

(139) FIG. **13** is a message sequence chart **1300** illustrating one implementation of a data processing system executing an application, where a Kernel Resource Manager (KRM) provides a routing table for the resource allocation for reconfigurable processors based on link bandwidth and link latency. A runtime logic **1310** compares a target link latency and target link bandwidth specification to rated link latency and rated link bandwidth values. Based on rated system dataflow resources available to meet target specifications, the runtime logic **1310** allocates pairs of reconfigurable processors and corresponding links for generating the configuration files used to load and execute an application.

(140) The Kernel Resource Manager (KRM) manages reconfigurable processors and associated hardware resources. It also has the view of the hardware topology that includes connectivity and proximity amongst hardware resources. When an application specifies hardware resources, the KRM examines the application against available resources and selects reconfigurable processors, device memory and links (e.g., PCIe lanes/ways) that maximize performance and minimize cost. The KRM provides a bandwidth and latency awareness for hardware resource partitioning and application scheduling.

(141) At operation one, an execution file comprising configuration files that define an application is received by the runtime logic **1310**. At operation two, the specification defining the application is processed by the runtime logic **1310**, where the execution file provides a configuration of virtual reconfigurable processors and virtual links. At operation three, the runtime logic **1310** processes the configuration encoded in a first virtual routing table that specifies a set of virtual links for each pair of virtual reconfigurable processors in the first plurality of virtual reconfigurable processors. At operation four, the specification is further encoded in the first virtual routing table that specifies a target link bandwidth of each virtual link in the set of virtual links. At operation five, the specification is further encoded in the virtual routing table that specifies a target link latency of each virtual link in the set of virtual links. At operation six, the Kernel Resource Manager **1320** creates a routing table that specifies a set of links for each pair of reconfigurable processors in the plurality of reconfigurable processors available to execute the first application. At operation seven, the runtime logic **1310** receives from the Kernel Resource Manager **1320** generated routing table specifying a rated link bandwidth of each link in the set of links. At operation eight, the runtime logic **1310** receives the from Kernel Resource Manager **1320** generated routing table that specifies a rated link latency of each link in the set of links. At operation nine, the runtime logic **1310** compares the rated link bandwidth and latency to target link bandwidth and latency of the virtual links defined by the specification. At operation ten, the runtime logic **1310** configures the allocated

pairs of reconfigurable processors and allocated links with the configuration data in the configuration files. At operation eleven, the runtime logic **1310** loads and executes the application using the configured pairs of reconfigurable processors and the configured links on a first processing node in the data center of FIG. 1.

(142) In some implementations, the specification includes metadata header files. In some implementations, operations seven and eight are reversed with operation eight preceding operation seven. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. 13. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. 13, operation eleven comprises streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(143) Having described the Kernel Resource Manager, we now return to executing and generating applications concurrently over the shared links for efficiently streaming data between processors on a same processing node and on different processing nodes of execution files disclosed by the technology disclosed.

(144) Executing First and Second Applications in Parallel

(145) FIG. 14 is a message sequence chart **1400** illustrating one implementation of a data processing system executing first and second applications using the disclosed resource allocation for reconfigurable processors based on link bandwidth and link latency. A runtime logic **1410** compares a target link latency and target link bandwidth specification to rated link latency and rated link bandwidth values. Based on rated system dataflow resources available to meet target specifications, the runtime logic **1410** allocates pairs of reconfigurable processors and corresponding links for generating the configuration files used to load and execute an application in parallel.

(146) At operation one, the runtime logic **1410** receives the first execution file comprising configuration files that define an application. At operation two, a second execution file comprising configuration files that define a second application is received by the runtime logic **1410**. At operation three, the configuration files, including the first and second configurations of virtual reconfigurable processors and virtual links defining each application, are processed by the runtime logic **1410**. At operation four, the runtime logic **1410** compares the rated link bandwidth to target link bandwidth of the virtual links defined by the first and second specifications. At operation five, the runtime logic **1410** compares the rated link latency to target link latency of the first and second virtual links specification. At operation six, the runtime logic **1410** configures and allocates pairs of reconfigurable processors and links according to the first and second specifications. At operation seven, the runtime logic **1410** loads reconfigurable processors (e.g., RPs **142**) with configuration files and executes the first and second applications. At operation eight, the reconfigurable processors process and stream first application data. At operation nine, the reconfigurable processors process and stream second application data. The streaming of data is on a first processing node in the data center of FIG. 1.

(147) In some implementations, the operations four and five are reversed with operation five preceding operation four. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. 14. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. 14, operation eleven comprises streaming network packets between reconfigurable processors (e.g., RPs **142**) and a host processor (e.g., host **102**) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(148) The application may include data and functions. Examples of the data include scalar data (e.g., control bits) and vector data (e.g., vectors, tensors, arguments, commands). The data can further include weights, coefficients, image data, audio data, Natural Language Processing data,

control data (e.g., control tokens). Examples of functions in the plurality of functions include non-linearities like Rectified Linear Unit (ReLU) and its variants (e.g., leaky ReLU), hyperbolic tangent, sigmoid, and softmax, element-wise addition, matrix multiplication (e.g., General Matrix Multiply (GeMM)), layer normalization (e.g., batch normalization), loss functions like cross-entropy, and tensor shape modifiers like transpose.

(149) One skilled in the art will appreciate that, in FIG. 14, streaming network packets between reconfigurable processors (e.g., RPs 142) and a host processor (e.g., host 102) on a same processing node over local links (e.g., PCIe links) can use a protocol like Transmission Control Protocol (TCP). Other links include for example (e.g., Ethernet, InfiniBand (IB)) using protocols like RDMA over Converged Ethernet (ROCE), TCP, User Datagram Protocol (UDP), and Quick UDP Internet Connections (QUIC).

(150) FIG. 15 is a message sequence chart 1500 illustrating one implementation of a data processing system executing first and second applications, such that first and second virtual routing tables are used for resource allocating reconfigurable processors based on link bandwidth and link latency.

(151) At operation one, runtime logic 1510 receives a first execution file comprising configuration files that define an application. At operation two, the runtime logic 1510 processes the first specification defining the application, where the execution file provides a first virtual routing table that specifies a first plurality of virtual reconfigurable processors required to execute the first application, and target link bandwidths and target link latencies of virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors. At operation three, the runtime logic 1510 receives a second execution file comprising configuration files that define an application. At operation four, the runtime logic 1510 receives a second virtual routing table that specifies a second plurality of virtual reconfigurable processors required to execute the second application, and target link bandwidths and target link latencies of virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors. At operation five, the runtime logic 1510 compares the two specifications encoded in the first and second virtual routing tables that specify a target link latency and a target link bandwidth of each virtual link in the first and second sets of virtual links and to the rated link latencies and rated link bandwidths. At operation six, the runtime logic 1510 configures and allocates pairs of reconfigurable processors and links as defined by the specifications according to the first and second virtual routing tables that meet rated requirements. At operation seven, the runtime logic 1510 loads and executes the first and second applications using the configured pairs of reconfigurable processors and the configured links on a first processing node in the data center of FIG. 1.

(152) In some implementations, operations seven and eight are reversed with operation eight preceding operation seven. Other implementations may perform the operations in different orders and/or with different, fewer, or additional operations than the ones illustrated in FIG. 15. Multiple operations can be combined in some implementations. One skilled in the art will appreciate that, in FIG. 15, operation eleven comprises streaming network packets between reconfigurable processors (e.g., RPs 142) and a host processor (e.g., host 102) on a same processing node over local links (e.g., PCIe links) using a protocol like Transmission Control Protocol (TCP).

(153) Example Reconfigurable Processor

(154) FIG. 16 is a diagram illustrating a system 1600 including a host 1620, a memory 1640, and a reconfigurable data processor 1610 in which a computation unit as described herein is deployed by hardware or by configuration of reconfigurable components and configured with the virtualization logic 1697. As shown in the example of FIG. 16, the reconfigurable data processor 1610 includes an array 1690 of configurable units and a configuration load/unload controller 1695.

(155) The virtualization logic 1697 can include resources that support or enable simultaneous execution of multiple, unrelated application graphs (or related ones) in an array of configurable

units on one die or one multi-chip modules. In the illustration, a first application graph is implemented in virtual machine VM1 in a particular set **1698** of configurable units, and a second application graph is implemented in virtual machine VM2 in another set **1699** of configurable units.

(156) Configurable units in an array **1690** of configurable units are further described in reference to FIGS. **19** and **20** and configured with the virtualization logic **1697**. Configurable units can include, or can have units configured to implement, a computation unit or computation units, as described herein.

(157) The reconfigurable data processor **1610** includes an external I/O interface **1630** connected to the host **1620** by line **1625**, and an external I/O interface **1650** connected to the memory **1640** by line **1645**. The I/O interfaces **1630**, **1650** connect via a bus system **1615** to the array **1690** of configurable units and to the configuration load/unload controller **1695**. The bus system **1615** may have a bus width of carrying one chunk of data, which can be for this example 128 bits (references to 128 bits throughout can be considered as an example chunk size more generally).

(158) To configure configurable units in the array **1690** of configurable units with a configuration file, the host **1620** can send the configuration file to the memory **1640** via the I/O interface **1630**, the bus system **1615**, and the I/O interface **1650** in the reconfigurable data processor **1610**. The configuration file can be loaded in many ways, as suits a particular architecture, including in data paths outside the reconfigurable configurable processor **1610**. The configuration file can be retrieved from the memory **1640** via the I/O interface **1650**. Chunks of the configuration file can then be sent in a distribution sequence to configurable units in the array **1690** of configurable units in the reconfigurable data processor **1610**.

(159) An external clock generator **1670** or other clock line sources can provide a clock line **1675** or clock lines to elements in the reconfigurable data processor **1610**, including the array **1690** of configurable units, and the bus system **1615**, and the external data I/O interfaces **1630** and **1650**. The bus system **1615** can communicate data at a processor clock rate via a clock line **1675** or clock lines.

(160) FIG. **17** is a simplified block diagram of components of a CGRA (Coarse-Grained Reconfigurable Architecture) processor. In this example, the CGRA processor has 2 tiles (Tile1, Tile2). The tile comprises an array of configurable units connected to a bus system, including array level networks in this example. An array of configurable units (e.g., **1690**, FIG. **16**) in the tile includes computation units in hardware or by configuration of reconfigurable components, which are configured with the virtualization logic **1697**. The bus system includes a top-level network connecting the tiles to external I/O interface **1705** (or any number of interfaces). In other embodiments, different bus system configurations may be utilized. The configurable units in each tile are nodes on the array level network in this embodiment.

(161) Each of the tiles has 4 AGCUs (Address Generation and Coalescing Units) (e.g., MAGCU1, AGCU17, AGCU13, AGCU14). The AGCUs are nodes on the top-level network and nodes on the array level networks and include resources for routing data among nodes on the top-level network and nodes on the array level network in each tile.

(162) Nodes on the top-level network in this example include one or more external I/So, including interface **1705**. The interfaces to external devices include resources for routing data among nodes on the top-level network and external devices, such as high-capacity memory, host processors, other CGRA processors, FPGA devices and so on, that are connected to the interfaces.

(163) One of the AGCUs in a tile is configured in this example to be a master AGCU, which includes an array configuration load/unload controller for the tile. In other embodiments, more than one array configuration load/unload controller can be implemented, and one array configuration load/unload controller may be implemented by logic distributed among more than one AGCU.

(164) The MAGCU1 includes a configuration load/unload controller for Tile1, and MAGCU2 includes a configuration load/unload controller for Tile2. In other embodiments, a configuration

load/unload controller can be designed for loading and unloading configuration of more than one tile. In other embodiments, more than one configuration controller can be designed for configuration of a single tile. Also, the configuration load/unload controller can be implemented in other portions of the system, including as a stand-alone node on the top-level network and the array level network or networks.

(165) The top-level network is constructed using top-level switches (**1711-1716**) connecting to each other as well as to other nodes on the top-level network, including the AGCUs, and I/O interface **1705**. The top-level network includes links (e.g., **L11**, **L9**, **L21**, **L22**) connecting the top-level switches. Data travels in packets between the top-level switches on the links, and from the switches to the nodes on the network connected to the switches. For example, top-level switches **1711** and **1712** are connected by a link **L11**, top-level switches **1714** and **1715** are connected by a link **L9**, top-level switches **1711** and **1714** are connected by a link **L13**, top-level switches **1715** and **1712** are connected by a link **L14**, top-level switches **1713** and **1716** are connected by a link **L15**, and top-level switches **1712** and **1713** are connected by a link **L21**. The links can include one or more buses and supporting control lines, including for example a chunk-wide bus (vector bus). For example, the top-level network can include data, request and response channels operable in coordination for transfer of data in a manner analogous to an AXI compatible protocol. Sec, AMBA® AXI and ACE Protocol Specification, ARM, 2017.

(166) Top-level switches can be connected to AGCUs. For example, top-level switches **1711**, **1712**, **1714**, and **1715** are connected to MAGCU1, AGCU17, AGCU14 and AGCU15 in the tile Tile1, respectively. Top-level switches **1712**, **1713**, **1715**, and **1716** are connected to MAGCU2, AGCU22, AGCU23 and AGCU24 in the tile Tile2, respectively.

(167) Top-level switches can be connected to one or more external I/O interfaces (e.g., interface **1705**).

(168) FIG. **18A** is a simplified diagram of a tile and an array level network usable in the configuration of FIG. **17**, where the configurable units in the array are nodes on the array level network and are configurable to implement a Look-Up Table with input offsetting. In this example, the array of configurable units **1800** includes a plurality of types of configurable units, which are configured with the virtualization logic **1697**. The types of configurable units in this example include Pattern Compute Units (PCU), Pattern Memory Units (PMU), Switch units(S), and Address Generation and Coalescing Units (each including two address generators AG and a shared CU). For an example of the functions of these types of configurable units, see, Prabhakar et al., “Plasticine: A Reconfigurable Architecture For Parallel Patterns,” ISCA '17, Jun. 24-28, 2017, Toronto, ON, Canada, which is incorporated by reference as if fully set forth herein. In this example, the PCUs (e.g., **1842**) and PMUs (e.g., **1843**) in the array of configurable units **1800** can include resources configurable for embodiment of a computation unit, an example configuration of which is described herein (FIGS. **18A**, **18B**, and **19**). Each of these configurable units contains a configuration store comprising a set of registers or flip-flops that represent either the setup or the sequence to run a program, and can include the number of nested loops, the limits of each loop iterator, the routes and/or instructions to be executed for each stage including stages, the source of the operands, and the network parameters for the input and output interfaces. The configuration file can include entries of lookup tables as described herein.

(169) Additionally, each of these configurable units contains a configuration store comprising a set of registers or flip-flops that store status usable to track progress in nested loops or otherwise. A configuration file in the configuration store contains a bit-stream representing the initial configuration, or starting state, of each of the components that execute the program. This bit-stream is referred to as a bit file. Program load is the process of setting up the configuration stores in the array of configurable units based on the contents of the bit file to allow the components to execute a program (i.e., a machine), including programs that utilize the virtualization logic **1697**. Program Load may also require the load of all PMU memories.

(170) The array level network includes links interconnecting configurable units in the array. The links in the array level network include one or more and, in this case, three kinds of physical buses: a chunk-level vector bus (e.g., 128 bits of data), a word-level scalar bus (e.g., 32 bits of data), and a multiple bit-level control bus. For instance, interconnect **1821** between switch units **1811** and **1812** includes a vector bus interconnect with a vector bus width of 128 bits, a scalar bus interconnect with a scalar bus width of 32 bits, and a control bus interconnect.

(171) The three kinds of physical buses differ in the granularity of data being transferred. In one embodiment, the vector bus can carry a chunk that includes 16-Bytes (=128 bits) of data as its payload. The scalar bus can have a 32-bit payload and carry scalar operands or control information. In some machines implemented using this system, data can be represented using floating point data formats, including standard or non-standard formats. Example formats include FP32 and BF16, among others. It can be understood that the number of data values carried on the scalar and vector buses is a function of the encoding format of the data values, with FP32 utilizing 32 bits per value and BF16 using 16 bits per value.

(172) The control bus can carry control handshakes such as tokens and other lines. The vector and scalar buses can be packet switched, including headers that indicate a destination of each packet and other information such as sequence numbers that can be used to reassemble a file when the packets are received out of order. Each packet header can contain a destination identifier that identifies the geographical coordinates of the destination Switch unit (e.g., the row and column in the array), and an interface identifier that identifies the interface on the destination switch (e.g., North, South, East, West, etc.) used to reach the destination unit. The control network can be circuit switched based on timing circuits in the device, for example. The configuration load/unload controller can generate a header for each chunk of configuration data of 128 bits. The header is transmitted on a header bus to each configurable unit in the array of configurable unit.

(173) In one example, a chunk of data of 128 bits is transmitted on the vector bus that provides the chunk as vector inputs to a configurable unit. The vector bus can include 128 payload lines, and a set of header lines. The header can include a sequence ID for each chunk, which can include: A bit to indicate if the chunk is scratchpad memory or configuration store data. Bits that form a chunk number. Bits that indicate a column identifier. Bits that indicate a row identifier. Bits that indicate a component identifier.

(174) For a load operation, the configuration load controller can send the number N of chunks to a configurable unit in order from N-1 to 0. For this example, the 6 chunks are sent out in most-significant-bit-first order of Chunk 5.fwdarw.Chunk 4.fwdarw.Chunk 3.fwdarw.Chunk 2.fwdarw.Chunk 1.fwdarw.Chunk 0. (Note that this most-significant-bit-first order results in Chunk 5 being distributed in round 0 of the distribution sequence from the array configuration load controller.) For an unload operation, the configuration unload controller can write the unload data out of order to the memory. For both load and unload operations, the shifting in the configuration serial chains in a configuration data store in a configurable unit is from LSB (Least-Significant-Bit) to MSB (Most-Significant-Bit), or MSB out first.

(175) FIG. **18B** illustrates an example Switch unit connecting elements in an array level network. As shown in the example of FIG. **18B**, a Switch unit can have 8 interfaces. The North, South, East and West interfaces of a switch unit are used for connections between switch units. The Northeast, Southeast, Northwest, and Southwest interfaces of a switch unit are each used to make connections to PCU or PMU instances. A set of 2 Switch units in each tile quadrant have connections to an Address Generation and Coalescing Unit (AGCU) that include multiple Address Generation (AG) units and a Coalescing Unit (CU) connected to the multiple Address Generation units. The Coalescing Unit (CU) arbitrates between the AGs and processes memory requests. Each of the 8 interfaces of a switch unit can include a vector interface, a scalar interface, and a control interface to communicate with the vector network, the scalar network, and the control network.

(176) During execution of a machine after configuration, data can be sent via one or more unit

switches and one or more links between the unit switches to the configurable units using the vector bus and vector interface(s) of the one or more Switch units on the array level network.

(177) In embodiments described herein, a configuration file or bit file, before configuration of the tile, can be sent from the configuration load controller using the same vector bus, via one or more unit switches and one or more links between the unit switches to the configurable unit using the vector bus and vector interface(s) of the one or more switch units on the array level network. For instance, a chunk of configuration data in a unit file particular to a configurable unit PMU **1841** can be sent from the configuration load/unload controller **1801** to the PMU **1841**, via a link **1820** between the configuration load/unload controller **1801** and the West (W) vector interface of the switch unit **1811**, the switch unit **1811**, and a link **1831** between the Southeast (SE) vector interface of the switch unit **1811** and the PMU **1841**.

(178) In this example, one of the AGCUs is configured to be a master AGCU, which includes a configuration load/unload controller (e.g., **1801**). The master AGCU implements a register through which the host (**1620**, FIG. **16**) can send commands via the bus system to the master AGCU. The master AGCU controls operations on an array of configurable units in a tile and implements a program control state machine to track the state of the tile based on the commands it receives from the host through writes to the register. For every state transition, the master AGCU issues commands to all components on the tile over a daisy-chained command bus (FIG. **16**). The commands include a program reset command to reset configurable units in an array of configurable units in a tile, and a program load command to load a configuration file to the configurable units.

(179) The configuration load controller in the master AGCU is responsible for reading the configuration file from the memory and sending the configuration data to every configurable unit of the tile. The master AGCU can read the configuration file from the memory at preferably the maximum throughput of the top-level network. The data read from memory are transmitted by the master AGCU over the vector interface on the array level network to the corresponding configurable unit according to a distribution sequence described herein.

(180) In one embodiment, in a way that can reduce the wiring requirements within a configurable unit, configuration and status registers holding unit files to be loaded in a configuration load process, or unloaded in a configuration unload process, in a component are connected in a serial chain and can be loaded through a process of shifting bits through the serial chain. In some embodiments, there may be more than one serial chain arranged in parallel or in series. When a configurable unit receives the for example 128 bits of configuration data from the master AGCU in one bus cycle, the configurable unit shifts this data through its serial chain at the rate of 1 bit per cycle, where shifter cycles can run at the same rate as the bus cycle. It will take 128 shifter cycles for a configurable unit to load 128 configuration bits with the 128 bits of data received over the vector interface. The 128 bits of configuration data are referred to as a chunk. A configurable unit can require multiple chunks of data to load all its configuration bits.

(181) The configurable units interface with the memory through multiple I/O interfaces (**1650**, FIG. **16**). Each of the memory interfaces can be accessed using several AGCUs. Each AGCU contains a reconfigurable scalar data path to generate requests for the off-chip memory. Each AGCU contains FIFOs (First-In-First-Out buffers for organizing data) to buffer outgoing commands, data, and incoming responses from the off-chip memory.

(182) FIG. **19** is a block diagram illustrating an example configurable unit **1900**, such as a Pattern Compute Unit (PCU), which is configured with the virtualization logic **1697**. A configurable unit can interface with the scalar, vector, and control buses, in this example using three corresponding sets of Inputs and Outputs (IO): scalar inputs/outputs, vector inputs/outputs, and control inputs/outputs. Scalar IOs can be used to communicate single words of data (e.g., 32 bits). Vector IOs can be used to communicate chunks of data (e.g., 128 bits), in cases such as receiving configuration data in a unit configuration load process and transmitting and receiving data during operation after configuration across a long pipeline between multiple PCUs. Control IOs can be

used to communicate signals on control lines such as the start or end of execution of a configurable unit. Control inputs are received by control block **1990**, and control outputs are provided by the control block **1990**.

(183) Each vector input is buffered in this example using a vector FIFO in a vector FIFO block **1960** which can include one or more vector FIFOs. Likewise, in this example, each scalar input is buffered using a scalar FIFO **1970**. Using input FIFOs decouples timing between data producers and consumers and simplifies inter-configurable-unit control logic by making it robust to input delay mismatches.

(184) A configurable unit includes multiple reconfigurable data paths in block **1980**. A data path in a configurable unit can be organized as a multi-stage (Stage 1 . . . Stage N), reconfigurable SIMD (Single Instruction, Multiple Data) pipeline. The chunks of data pushed into the configuration serial chain in a configurable unit include configuration data for each stage of each data path in the configurable unit. The configuration serial chain in the configuration data store **1920** is connected to the multiple data paths in block **1980** via lines **1921**.

(185) A configurable data path organized as a multi-stage pipeline can include multiple functional units (e.g., **1981**, **1982**, **1983**, **1984**, **1985**, **1986**) at respective stages. A computation unit or parts of a computation unit can be implemented in multiple functional units at respective stages in a multi-stage pipeline or in multiple multi-stage pipelines, as described in FIG. **15**. In the example as shown in FIG. **15**, a circuit including the virtualization logic **1697** can be implemented in multiple functional units and multiple memory units. Input registers in functional units can register inputs from scalar FIFOs **1970** or Vector FIFOs **1960** or from previous stages in a multi-stage pipeline. A functional unit at a stage in a multi-stage pipeline can execute a function, e.g., logical shift, an arithmetic function, comparison, a logical operation, etc., and generate an output.

(186) Configurable units in the array of configurable units include configuration data stores **1920** (e.g., serial chains) to store unit files comprising a plurality of chunks (or sub-files of other sizes) of configuration data particular to the corresponding configurable units. Configurable units in the array of configurable units each include unit configuration load/unload process logic **1940** connected to the configuration data store **1920** via line **1922**, to execute a unit configuration load process. The unit configuration load process includes receiving, via the bus system (e.g., the vector inputs), chunks of a unit file particular to the configurable unit and loading the received chunks into the configuration data store **1920** of the configurable unit. The unit file loaded into the configuration data store **1920** can include configuration data, including opcodes and routing configuration, for circuits (e.g., module) implementing the virtualization logic **1697** in multiple functional units and multiple memory units, as described herein.

(187) The configuration data stores in configurable units in the plurality of configurable units in this example comprise serial chains of latches, where the latches store bits that control configuration of the resources in the configurable unit. A serial chain in a configuration data store can include a shift register chain for configuration data and a second shift register chain for state information and counter values connected in series.

(188) Input configuration data **1910** can be provided to a vector FIFO as vector inputs, and then be transferred to the configuration data store **1920**. Output configuration data **1930** can be unloaded from the configuration data store **1920** using the vector outputs.

(189) The CGRA uses a daisy-chained completion bus to indicate when a load/unload command has been completed. The master AGCU transmits the program load and unload commands to configurable units in the array of configurable units over a daisy-chained command bus. As shown in the example of FIG. **19**, a control block **1990**, a daisy-chained completion bus **1991** and a daisy-chained command bus **1992** are connected to daisy-chain logic **1993**, which communicates with the unit configuration load logic **1940**. The daisy-chain logic **1993** can include load complete status logic, as described below. The daisy-chained completion bus is further described below. Other topologies for the command and completion buses are clearly possible but not described here.

(190) FIG. 20 is a block diagram illustrating an example configurable unit **2000**, such as a Pattern Memory Unit (PMU), which is configured with the virtualization logic **1697** (i.e., the ready-to-read credit counter, the write credit counter, and the flow control logic for operating them). A PMU can contain scratchpad memory **2030** coupled with a reconfigurable scalar data path **2020** intended for address calculation (RA, WA) and control (WE, RE) of the scratchpad memory **2030**, along with the bus interfaces used in the PCU (FIG. 19).

(191) The bus interfaces can include scalar inputs, vector inputs, scalar outputs and vector outputs, usable to provide Write Data WD. The data path can be organized as a multi-stage reconfigurable pipeline, including stages of Functional Units FUs and associated Pipeline Registers PRs that register inputs and outputs of the functional units. PMUs can be used to store distributed on-chip memory throughout the array of reconfigurable units.

(192) A scratchpad is built with multiple SRAM banks (e.g., **2031-2034**). Banking buffering logic **2035** for the SRAM banks in the scratchpad can be configured to operate in several banking modes to support various access patterns. A computation unit as described herein can include a lookup table stored in the scratchpad memory **2030**, from a configuration file or from other sources. In a computation unit as described herein, the scalar data path can translate a section of a raw input value I for addressing lookup tables implementing a function $f(I)$, into the addressing format utilized by the SRAM scratchpad memory **2030**, adding appropriate offsets and so on, to read the entries of the lookup table stored in the scratchpad memory **2030** using the sections of the input value I. Each PMU can include write address calculation logic and read address calculation logic **2020** that provide Write Address WA, Write Enable WE, Read Address RA and Read Enable RE to the banking buffering logic **2035**. Based on the state of the local FIFOs **2011** and **2012** and external control inputs, the control block **2020** can be configured to trigger the Write Address computation, Read Address computation, or both, by enabling the appropriate counters **2016**. A programmable counter chain (Control Inputs, Control Outputs) and control block **2020** can trigger PMU execution.

(193) This is one simplified example of a configuration of a configurable processor for implementing a computation unit as described herein. The configurable processor can be configured in other ways to implement a computation unit. Other types of configurable processors can implement the computation unit in other ways. Also, the computation unit can be implemented using dedicated logic in some examples, or a combination of dedicated logic and instruction-controlled processors.

OTHER IMPLEMENTATIONS

(194) A first example of accelerated deep learning is using a deep learning accelerator to train a neural network. A second example of accelerated deep learning is using a deep learning accelerator to operate a trained neural network to perform inferences. A third example of accelerated deep learning is using a deep learning accelerator to train a neural network and subsequently perform inference with any one or more of the trained neural networks, information from same, and a variant of same.

(195) Examples of neural networks include Fully Connected Neural Networks (FCNNs), Recurrent Neural Networks (RNNs), Convolutional Neural Networks (CNNs), Long Short-Term Memory (LSTM) networks, autoencoders, deep belief networks, and Generative Adversarial Networks (GANs).

(196) An example of training a neural network is determining one or more weights associated with the neural network, such as by hardware acceleration via a deep learning accelerator. An example of making an inference is using a trained neural network to compute results by processing input data based on weights associated with the trained neural network. As used herein, the term 'weight' is an example of a 'parameter' as used in various forms of neural network processing. For example, some neural network learning is directed to determining parameters that are then usable for performing neural network inferences using the parameters.

(197) A neural network processes data according to a dataflow graph comprising layers of neurons. Stimuli (e.g., input data) are received by an input layer of neurons and the computed results of the dataflow graph (e.g., output data) are provided by an output layer of neurons. Example layers of neurons include input layers, output layers, rectified linear unit layers, fully connected layers, recurrent layers, long short-term memory layers, convolutional layers, kernel layers, dropout layers, and pooling layers. A neural network is conditionally and/or selectively trained, subject to hardware acceleration. After being trained, a neural network is conditionally and/or selectively used for inference, subject to hardware acceleration.

(198) An example of a deep learning accelerator (chip) is one or more relatively specialized hardware elements operating in conjunction with one or more software elements to train a neural network and/or perform inference with a neural network relatively more efficiently than using relatively less specialized hardware elements. Some implementations of the relatively specialized hardware elements include one or more hardware logic circuitry elements such as transistors, resistors, inductors, capacitors, wire interconnects, combinatorial logic (e.g., NAND, NOR) gates, latches, register files, memory arrays, tags for memory arrays, content-addressable memories, flash, ROM, DRAM, SRAM, Serializer/Deserializer (SerDes), I/O drivers, and the like, such as implemented via custom logic, synthesized logic, ASICs, and/or FPGAs. Some of the relatively less specialized hardware elements include conventional CPUs and conventional GPUs.

(199) An example of storage is one or more elements enabled to retain state information, e.g., any one or more of: a flip-flop, a latch or an array of latches, a register or an array of registers, a register file, a memory, a memory array, a magnetic storage device, an optical storage device, SRAM, DRAM, flash, and ROM. In various implementations storage is volatile (e.g., SRAM or DRAM) and/or non-volatile (e.g., flash or ROM).

(200) An example of an Integrated Circuit (IC) is a collection of circuitries implemented on one or more portions of semiconductor material, such as a single die or a plurality of dice. An example of 3D-stacking of dice is providing mechanical connectivity and/or electrical connectivity between the dice, e.g., in a dimension orthogonal to a major surface of the dice, to form a unit. The mechanical connectivity and/or the electrical connectivity are variously implemented, e.g., via one or more of solder balls, microbumps, and through-silicon vias. An example of 2.5D stacking of dice is providing mechanical connectivity and/or electrical connectivity between the dice via a common element (e.g., a silicon interposer) to form a unit, where the mechanical connectivity and/or electrical connectivity between each die and the common substrate is in a dimension orthogonal to a major surface of the die. The mechanical connectivity and/or the electrical connectivity are variously implemented, e.g., via one or more of solder balls, microbumps, and through-silicon vias. An example of an Application-Specific Integrated Circuit (ASIC) is an IC designed for a particular use.

(201) An example of a package is an element enabled to mechanically retain and/or contain one or more electronic circuits and/or to electrically interconnect one or more electronic circuits. Example electronic circuits are any one or more of one or more portions of semiconductor material, one or more dice, one or more interposers, and one or more substrates. Particular examples of packages include a BGA package and variants thereof. Some ICs comprise a package. An example of a substrate is an element to mechanically retain and/or electrically interconnect one or more dice and/or one or more packages. A particular example of a substrate is a PCB to, e.g., retain and interconnect packages. Another particular example of a substrate is a silicon interposer to, e.g., couple one or more 3D-stacked or 2.5-stacked dice. Another particular example of a substrate is a package, e.g., retaining a plurality of dice.

(202) The technology disclosed can be applied to other processors like Central Processing Units (CPUs), Graphics Processing Units (GPUs), Field Programmable Gate Arrays (FPGAs), Application-Specific Integrated Circuits (ASICs), Application Specific Instruction-set Processor (ASIP), and Digital Signal Processors (DSPs).

(203) The applications can also be considered graphs, application graphs, dataflow graphs, control flow graphs, data and control flow graphs, user applications, models, deep learning applications, deep neural networks, programs, program images, processes, jobs, and tasks.

(204) While the present invention is disclosed by reference to the preferred implementations and examples detailed above, it is to be understood that these examples are intended in an illustrative rather than in a limiting sense. It is contemplated that modifications and combinations will readily occur to those skilled in the art, which modifications and combinations will be within the spirit of the invention and the scope of the following clauses.

Claims

1. A data processing system, comprising: a pool of reconfigurable processors operatively coupled by links, the links having rated link bandwidths and rated link latencies; runtime logic operatively coupled to the pool of reconfigurable processors, the runtime logic configured to receive, for a first application, a first plurality of configuration files containing configuration data; a first configuration of a first plurality of virtual reconfigurable processors required to execute the first application; virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; and a first specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; using the runtime logic to allocate the reconfigurable processors in the plurality of reconfigurable processors to the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; using the runtime logic to allocate links between the reconfigurable processors to the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors based on a link bandwidth comparison that compares the target link bandwidths, specified by the first specification, against the rated link bandwidths, and a link latency comparison that compares the target link latencies, specified by the first specification, against the rated link latencies; and using the runtime logic to allocate the reconfigurable processors and the links with configuration data in the first plurality of configuration files, forming configured reconfigurable processors and configured links; and executing the first application using the configured reconfigurable processors and the configured links.
2. The data processing system of claim 1, wherein the runtime logic is further configured to receive, for a second application, a second plurality of configuration files that contain configuration data; a second configuration of a second plurality of virtual reconfigurable processors required to execute the second application; virtual links between virtual reconfigurable processors in the second plurality of virtual reconfigurable processors; and a second specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors.
3. The data processing system of claim 1, wherein the pool of reconfigurable processors comprise a cluster of processing nodes connected through a network.
4. The data processing system of claim 1, wherein one or more reconfigurable processors are host processors for controlling dataflow through a network.
5. The data processing system of claim 3, wherein the pool of reconfigurable processors is associated with a single processing node in a network.
6. The data processing system of claim 3, wherein the pool of reconfigurable processors is associated with multiple processing nodes in a network.
7. The data processing system of claim 3, where the allocated reconfigurable processors are on a same processing node.
8. The data processing system of claim 3, where the allocated reconfigurable processors are on a different processing node.

9. The data processing system of claim 1, wherein the pool of reconfigurable processors is dynamically scalable to meet the performance requirements of applications requesting execution.
 10. The system of claim 1, wherein the processor elements are respective arrays of configurable units.
 11. The system of claim 5, wherein the reconfigurable processors are pattern compute units (PCUs) and pattern memory units (PMUs).
 12. A computer-implemented method, comprising: providing a pool of reconfigurable processors operatively coupled by links, the links having rated link bandwidths and rated link latencies; providing runtime logic operatively coupled to the pool of reconfigurable processors, the runtime logic configured to receive, for a first application, a first plurality of configuration files containing configuration data; providing a first configuration of a first plurality of virtual reconfigurable processors required to execute the first application; providing virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; and providing a first specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; using the runtime logic, allocating reconfigurable processors in the plurality of reconfigurable processors to the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors; using the runtime logic, allocating links between the reconfigurable processors to the virtual links between the virtual reconfigurable processors in the first plurality of virtual reconfigurable processors based on a link bandwidth comparison that compares the target link bandwidths, specified by the first specification, against the rated link bandwidths, and a link latency comparison that compares the target link latencies, specified by the first specification, against the rated link latencies; using the runtime logic, allocating the reconfigurable processors and the links with configuration data in the first plurality of configuration files, forming configured reconfigurable processors and configured links; and executing the first application using the configured reconfigurable processors and the configured links.
 13. The computer-implemented method of claim 12, wherein the runtime logic is further configured to receive, for a second application, a second plurality of configuration files that contain configuration data; a second configuration of a second plurality of virtual reconfigurable processors required to execute the second application; virtual links between virtual reconfigurable processors in the second plurality of virtual reconfigurable processors; and a second specification of target link bandwidths and target link latencies of the virtual links between the virtual reconfigurable processors in the second plurality of virtual reconfigurable processors.
 14. The computer-implemented method of claim 12, wherein the bandwidth and latency are achievable bandwidth and achievable latency.
 15. The computer-implemented method of claim 12, wherein the pool of reconfigurable processors comprise a cluster of processing nodes connected through a network.
 16. The computer-implemented method of claim 12, wherein one or more reconfigurable processors are host processors for controlling dataflow through a network.
 17. The computer-implemented method of claim 15, wherein the pool of reconfigurable processors is a single processing node or multiple processing nodes coupled to a plurality of reconfigurable processors.
 18. The computer-implemented method of claim 15, wherein the allocated reconfigurable processors are on a same processing node.
 19. The computer-implemented method of claim 15, where the allocated reconfigurable processors are on a different processing node.
 20. The computer-implemented method of claim 12, wherein the pool of reconfigurable processors is dynamically scalable to meet the performance requirements of applications requesting execution.
-