

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12413535
Kind Code	B1
Date of Patent	September 09, 2025
Inventor(s)	Jain; Ajit Kumar et al.

Efficient scheduling using adaptive packing mechanism for network apparatuses

Abstract

A scheduler in a network device serves ports with data units from a plurality of queues. The scheduler implements a scheduling algorithm that is normally constrained to releasing data to a port no more frequently than at a default maximum service rate. However, when data units smaller than a certain size are at the heads of one or more data unit queues assigned to a port, the scheduler may temporarily increase the maximum service rate of that port. The increased service rate permits fuller realization of a port's maximum bandwidth when handling smaller data units. In some embodiments, increasing the service rate involves dequeuing more than one small data unit at a time, with the extra data units temporarily stored in a port FIFO. The scheduler adds a pseudo-port to its scheduling sequence to schedule release of data from the port FIFO, with otherwise minimal impact on the scheduling logic.

Inventors:	Jain; Ajit Kumar (Milpitas, CA), Alapati; Ashwin (San Jose, CA)
Applicant:	Innovium, Inc. (Santa Clara, CA)
Family ID:	88534561
Assignee:	Innovium, Inc. (Santa Clara, CA)
Appl. No.:	18/385342
Filed:	October 30, 2023

Related U.S. Application Data

continuation parent-doc US 17141130 20210104 US 11805066 child-doc US 18385342

Publication Classification

Int. Cl.: H04L47/62 (20220101); H04L45/00 (20220101); H04L47/52 (20220101); H04L49/90 (20220101)

U.S. Cl.:

CPC H04L47/6225 (20130101); H04L45/68 (20130101); H04L47/521 (20130101); H04L47/6215 (20130101); H04L49/90 (20130101);

Field of Classification Search

CPC: H04L (47/6225); H04L (47/6215); H04L (47/521); H04L (45/68); H04L (49/90)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
8588242	12/2012	Izenberg	710/52	H04L 47/527
11805066	12/2022	Jain	N/A	H04L 45/68
2006/0187949	12/2005	Seshan	370/229	H04L 49/901

OTHER PUBLICATIONS

United States Patent and Trademark Office, U.S. Appl. No. 17/141,130, Non-Final Office Action dated Feb. 17, 2023. cited by applicant

United States Patent and Trademark Office, U.S. Appl. No. 17/141,130, Notice of Allowance dated Jun. 12, 2023. cited by applicant

Primary Examiner: Liu; Siming

Attorney, Agent or Firm: Shield Intellectual Property PC

Background/Summary

TECHNICAL FIELD

(1) Embodiments relate generally to computer networking, and, more specifically, to techniques for improving a network device's throughput via enhancements to its scheduler(s).

BACKGROUND

(2) The approaches described in this section are approaches that could be pursued, but not necessarily approaches that have been previously conceived or pursued. Therefore, unless otherwise indicated, it should not be assumed that any of the approaches described in this section qualify as prior art merely by virtue of their inclusion in this section.

(3) A computer network is a set of computing components interconnected by communication links. Each computing component may be a separate computing device, such as, without limitation, a hub, switch, bridge, router, server, gateway, or personal computer, or a component thereof. Each computing component, or “network device,” is considered to be a node within the network. A communication link is a mechanism of connecting at least two nodes such that each node may transmit data to and receive data from the other node. Such data may be transmitted in the form of signals over transmission media such as, without limitation, electrical cables, optical cables, or wireless media.

(4) The structure and transmission of data between nodes is governed by a number of different protocols. There may be multiple layers of protocols, typically beginning with a lowest layer, such as a “physical” layer that governs the transmission and reception of raw bit streams as signals over a transmission medium. Each layer defines a data unit (the protocol data unit, or “PDU”), with multiple data units at one layer typically combining to form a single data unit in another. Additional examples of layers may include, for instance, a data link layer in which bits defined by a physical layer are combined to form a frame or cell, a network layer in which frames or cells defined by the data link layer are combined to form a packet, and a transport layer in which packets defined by the network layer are combined to form a TCP segment or UDP datagram. The Open Systems Interconnection model of communications describes these and other layers of communications. However, other models defining other ways of layering information may also be used. The Internet protocol suite, or “TCP/IP stack,” is one example of a common group of protocols that may be used together over multiple layers to communicate information. However, techniques described herein may have application to other protocols outside of the TCP/IP stack.

(5) A given node in a network may not necessarily have a link to each other node in the network, particularly in more complex networks. For example, in wired networks, each node may only have a limited number of physical ports into which cables may be plugged in to create links. Certain “terminal” nodes—often servers or end-user devices—may only have one or a handful of ports. Other nodes, such as switches, hubs, or routers, may have a great deal more ports, and typically are used to relay information between the terminal nodes. The arrangement of nodes and links in a network is said to be the topology of the network, and is typically visualized as a network graph or tree.

(6) A given node in the network may communicate with another node in the network by sending data units along one or more different paths through the network that lead to the other node, each path including any number of intermediate nodes. The transmission of data across a computing network typically involves sending units of data, such as packets, cells, or frames, along paths through intermediary networking devices, such as switches or routers, that direct or redirect each data unit towards a corresponding destination.

(7) While a data unit is passing through an intermediary networking device—a period of time that is conceptualized as a “visit” or “hop”—the device may perform any of a variety of actions, or processing steps, with the data unit. The exact set of actions taken will depend on a variety of characteristics of the data unit, such as metadata found in the header of the data unit, and in many cases the context or state of the network device. For example, address information specified by or otherwise associated with the data unit, such as a source address, a destination address, or path information, is typically used to determine how to handle a data unit (e.g. what actions to take with respect to the data unit). For instance, an Internet Protocol (“IP”) data packet may include a destination IP address field within the header of the IP data packet, based upon which a network device may determine one or more other networking devices, among a number of possible other networking devices, to forward the IP data packet to. The logic within a network device that controls the specific set of actions performed with respect to a given data unit is often referred to as “packet-switching” or “forwarding” logic.

(8) A network device may include any number of internal hardware and/or software components configured to handle the movement of data units between processing components within the device and, eventually, out of the device. It is desirable for these components to quickly determine where to send and/or store data for processing, and to expediently send and/or store that data to the appropriate destination once determined. Moreover, it is desirable for these components to handle network traffic in a manner that will optimally utilize available network resources throughout the network in which the device is situated.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present inventive subject matter is illustrated by way of example, and not by way of limitation, in the figures of the accompanying drawings and in which like reference numerals refer to similar elements and in which:
- (2) FIG. 1 is an illustrative view of various components of an example scheduler-based subsystem for queueing data units while those data units await transmission to downstream logic within a network device comprising a plurality of ports;
- (3) FIG. 2 illustrates a scheduler-based dequeue subsystem configured to utilize port FIFOs;
- (4) FIG. 3 is an illustrative view of components for a scheduler-based queueing subsystem, with respect to an individual port that utilizes a port FIFO;
- (5) FIG. 4 illustrates an example flow for enqueueing a data unit in a queue;
- (6) FIG. 5 illustrates an example flow for scheduling a dequeue operation;
- (7) FIG. 6 illustrates an example flow for increasing the service rate of a port for smaller data units;
- (8) FIG. 7 illustrates an example flow for implementing a dequeue operation in response to a request to dequeue a normal port that is coupled to a port FIFO;
- (9) FIG. 8 illustrates an example flow for implementing a dequeue operation in response to a request to dequeue a pseudo-port associated with a port FIFO;
- (10) FIG. 9 is an illustrative view of various aspects of an example networking system in which the techniques described herein may be practiced;
- (11) FIG. 10 is an illustrative view of various aspects of an example network device in which techniques described herein may be practiced;
- (12) FIG. 11 illustrates an example of a network device with multiple packet processing pipelines; and
- (13) FIG. 12 is a block diagram that illustrates an example computer system that may be utilized in implementing the above-described techniques.

DETAILED DESCRIPTION

(14) In the following description, for the purposes of explanation, numerous specific details are set forth in order to provide a thorough understanding of the present inventive subject matter. It will be apparent, however, that the present inventive subject matter may be practiced without these specific details. In other instances, well-known structures and devices are shown in block diagram form in order to avoid unnecessarily obscuring the present inventive subject matter.

(15) Embodiments are described herein according to the following outline: 1.0 General Overview 2.0. System Overview 2.1. Data Units 2.2. Ports 2.3. Queue Manager 2.4. Scheduler 2.5. Port FIFOs and Pseudo-Ports 2.6. Example Port FIFO Implementation for Single Port 2.7. Example Port FIFO Operation 2.8. Miscellaneous 3.0. Functional Overview 3.1. Enqueueing a Data Unit 3.2. Scheduling Dequeue Operations 3.3. Increasing Service Rate for Small Data Units 3.4. Dequeueing a Normal Port Coupled to a Port FIFO 3.5. Dequeueing a Pseudo Port for a Port FIFO 4.0. Example Embodiments 5.0. Example Network Device 5.1. Networks 5.2. Network Device 5.3. Ports 5.4. Packet Processors 5.5. Buffers 5.6. Queues 5.7. Traffic Management 5.8. Forwarding Logic 5.9. Multi-Pipeline Architecture 5.10. Miscellaneous 6.0. Implementation Mechanism—Hardware Overview 7.0. Extensions and Alternatives

1.0. General Overview

(16) Approaches, techniques, and mechanisms are disclosed for improving performance of switches or other network devices by increasing the efficiency with which a scheduler serves a plurality of ports with data units from a plurality of queues. The scheduler implements a scheduling algorithm that is normally constrained to scheduling the release of data to a port no more frequently than at a default maximum service rate (e.g. once every two clock cycles, once every four clock cycles, etc.).

However, when data units smaller than a certain size are at the heads of one or more queues of data units assigned to a port, the scheduler may temporarily increase the maximum service rate of that port. The maximum service rate may return to its default value, for instance, when a data unit larger than the certain size returns to the head of the queue.

(17) Among other advantages, the improved scheduler permits fuller realization of a port's maximum bandwidth when handling smaller packets or other data units. For example, a scheduler dealing with a stream of 64 byte packets destined for certain port on a device running at a certain clock speed might have to provide a new packet to the port at a frequency greater than every other clock cycle. This would mean the scheduler would have to be capable of scheduling packets for dispatch to the port on back-to-back clock cycles at least some of the time to realize the full capacity of the port. However, it is typically not desirable for a scheduler to do so. Serving a port at high frequencies adds technical complexity, such as requiring extremely fast updating and communication of control information concerning the availability of data for the port to the scheduler (potentially prohibitively so). Moreover, permitting a scheduler to serve a port at high frequencies may aggravate burst behavior within the device. That is, if data units are very large instead of very small, the scheduler may end up sending more data downstream than can be handled at the port (or another downstream component), meaning that the downstream component would either have to drop data units and/or that the downstream component may send instructions upstream to stop transmitting to the port for a time. Such an instruction is typically undesirable, as it may cause upstream components to pause transmission for longer than necessary, result in dropped data units, and/or have other detrimental effects.

(18) Hence, according to an embodiment, the scheduler is permitted to service a port more frequently (e.g. back-to-back) only when serving data units smaller than a certain size, and reverts to a less frequent maximum service rate when serving data units larger than that size. The size may vary depending on the embodiment as well as the context. For instance, in an embodiment, the data units are transport data units ("TDU") that are typically of a fixed maximum size. Any data unit smaller than the fixed maximum size may activate the scheduler's ability to serve data units more frequently. In an embodiment, the scheduler's behavior is activated only when transmitting data units that are considered to be end-of-packet ("EOP") data units followed by start-of-packet ("SOP") data units of less than a certain size. In an embodiment, a data unit must also be small enough to be transmitted on a bus of a fixed size concurrently with the previous or next data unit to be transmitted to the port. For instance, if the bus size is 512 bytes, and a first data unit is 256 bytes in size, if the next data unit is no greater than 256 bytes in size, then the scheduler's ability to serve data units more frequently may be activated.

(19) According to an embodiment, the ability to schedule a port for service more frequently than the default maximum service rate is enabled by the introduction of a combination of features. These features include rate adjustment logic configured to recognize when multiple data units are ready to be dequeued, dequeue logic capable of dequeuing these multiple data units simultaneously, transmission logic configured to send the first of these multiple data units downstream while holding the other(s) temporarily in a first-in-first-out buffer for the port (referred to herein as the "port FIFO"), and logic at the scheduler that associates the port FIFO with a pseudo-port that, when scheduled, causes a data unit to be released from the FIFO. Among other advantages, such a combination may, in some embodiments, require minimal changes to the scheduler and/or other components of a network device compared to certain alternative implementations.

(20) According to an embodiment, the foregoing techniques may be implemented on a port-by-port basis. For example, a device manufacturer may advertise certain ports as having the ability to increase service rates for small packets, whereas other ports might not. If a higher service rate is enabled by means of a port FIFO, a port FIFO may be provided for some ports, but, to reduce expense, not for others. For instance, it may be desirable to include port FIFOs to permit higher service rates for high speed ports (e.g. 400 Gb/s, 800 Gb/s, etc.), but not for lower speed ports (e.g.

100 Gb/s). In an embodiment, the foregoing functionality may be enabled programmatically by a user or device manufacturer.

2.0. System Overview

(21) FIG. 1 is an illustrative view of various components of an example scheduler-based subsystem **100** for queueing data units while those data units await transmission to downstream logic within a network device comprising a plurality of ports **190**, according to an embodiment. Subsystem **100** may be a subsystem within a switch or other network device, as described elsewhere herein, or any other computing device. In an embodiment, the various components of subsystem **100** described below are hardware-based logic units within one or more application-specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), or other integrated circuit(s). In other embodiments, some or all of the components may be software-based logic implemented by one or more processors executing instructions stored in one or more computer-readable media.

(22) Subsystem **100** may be utilized in any component of a device that needs to queue data destined for or otherwise associated with different ports **190** (or other constructs). For instance, subsystem **100** may be utilized in a network device to manage the flow of data units through queues within a traffic manager or ingress arbiter to downstream logic such as egress ports, egress processors, ingress processors, etc. The example components of subsystem **100** depicted in FIG. 1 are now described.

(23) 2.1. Data Units

(24) Data is communicated to and from subsystem **100** as series of discrete units referred to as data units **105**. Data units **105** may be any type of data unit. In an embodiment, a data unit **105** is a set of data no larger than a fixed maximum size (e.g. a fixed number of bits or bytes) for transmission and/or storage purposes within a device. This fixed maximum size may correspond to, for instance, the width(s) of certain busses or interconnects within the device that implements subsystem **100**, the maximum amount of data that may be processed by certain components within the device at a time, the maximum size of a row or entry in a buffer that stores each data unit **105**, and/or other constraints.

(25) In an embodiment, each data unit **105** is a portion of a larger data unit. To differentiate between these two types of data units, the larger data units may be referred to as Protocol Data Units (“PDUs”) while data units **105** may be referred to as Transport Data Units (“TDUs”). Some PDUs may be of a significantly larger size than a TDU. For instance, a PDU may be divided into tens, hundreds, or even greater numbers of constituent TDUs. However, smaller PDUs may also exist, including PDUs that are smaller than the fixed maximum size of a TDU, and are thus referred to as single-TDU PDUs.

(26) A PDU may have been divided into TDUs at some point prior to arrival at subsystem **100**, such as by ingress logic upon receipt of the PDU at the implementing device, or even by another device that sent the PDU. TDUs may be accompanied by metadata, such as an internal header or sideband control information, that identifies the PDU to which the TDU belongs. TDUs belonging to the same PDU are linked together, and will typically arrive at and depart from the subsystem **100** in their order within the PDU. The first TDU within a PDU is known as the start-of-PDU or start-of-packet (“SOP”) TDU, while the last TDU within the PDU is known as the end-of-PDU or end-of-packet (“EOP”) TDU. All TDUs between the SOP and EOP are referred to as middle-of-packet (“MOP”) TDUs. In an embodiment, all TDUs except for EOP TDUs are of the same fixed maximum size. In an embodiment, a SOP TDU that is also an EOP TDU (that is, a single-TDU PDU), may be referred to as an SEOP TDU.

(27) In an embodiment, each PDU is an Internet Protocol (“IP”) packet, while each TDU is a cell or frame of an IP packet. However, in other embodiments, PDUs and TDUs may conform to other structures, such as segments, datagrams, messages, etc. In yet other embodiments, each data unit **105** may instead be a PDU.

(28) Data units **105** may be stored temporarily in a buffer (not depicted). They may have been

written to the buffer by any suitable component of the implementing device. For instance, the data units **105** may have been received by a traffic manager, ingress arbiter, or other component with limited resources for handling data units **105**. Such a component may have placed a data unit **105** in the buffer until it is ready to process and/or transmit the data unit **105**. In some embodiments, a TDU may itself be subdivided into Storage Data Units (“SDUs”) for storage purposes within the buffer, but is combined back together for transmission purposes.

(29) 2.2. Ports

(30) Each data unit **105** is to be transmitted to one or more ports **190**. Depending on where subsystem **100** is deployed within an implementing device, subsystem **100** may transmit a data unit **105** directly to a port **190**, or subsystem **100** may transmit a data unit to downstream logic comprising one or more components that further processes the data unit **105** and then forward the data unit **105** to its destination port(s) **190**. In the former case, subsystem **100** may thus have separate egress paths for each port **190**, while in the latter case, subsystem **100** may or may not have separate egress paths for each port **190**, depending on the configuration of the downstream logic.

(31) Metadata associated with a data unit **105** (e.g. header information, control information, etc.) indicates the port(s) **190** to which a data unit **105** is to be transmitted. In an embodiment, such metadata may have been generated by the implementing device in response to analyzing the data unit **105**, or other data units **105** belonging to the same PDU. For instance, forwarding logic within the implementing device may have analyzed the SOP TDU of a PDU upon arrival at the implementing device. By comparing destination address(es) and/or other header information within the SOP TDU to forwarding instructions within a forwarding table, the forwarding logic may have identified one or more ports **190** to send the PDU to. The forwarding logic may have generated control information for the PDU that identifies these port(s) **190**, and then linked this control information to each TDU in the PDU (e.g. by inserting it into each TDU, prepending each TDU with the control information, linking each TDU with a PDU identifier that in turn is linked to the control information, etc.).

(32) 2.3. Queue Manager

(33) Subsystem **100** includes a queue manager **110** that maintains queues **140** of data units **105** that are ready for transmission to ports **190**. A queue **140** is a sequence of zero or more of the data units **105**, which indicates an order for processing those data units **105**. The first data unit **105** in the sequence is said to be at the “head” of the queue **140**, while the last data unit **105** is said to be at the tail. When the first data unit **105** is ready to be processed, it is dequeued, or “popped,” from the queue **140**, and the next data unit **105** in the sequence becomes the head. In this manner, a data unit **105** moves from the tail of the queue **140** to the head of the queue **140** over time, and is eventually dequeued.

(34) The size of a queue **140** will vary over time. In an embodiment, there may be a maximum queue size. A queue **140** may be described by linking data, such as a linked list of data unit identifying information, an array of data unit identifying information, or any other suitable structure. In yet other embodiments, a queue **140** may be a distinct portion of a buffer in which its constituent data units **105** are stored in the corresponding sequence.

ENQUEUE LOGIC

(35) Queue manager **110** includes enqueue logic **120** configured to receive a data unit **105**, or information identifying a data unit **105**, from upstream logic within the implementing device. In the former case, if a data unit **105** is received directly by enqueue logic **120**, queue manager **110** may interact with buffer management logic to store the data unit **105** in a suitable location. In the latter case, the information may be, for instance, metadata for the data unit **105**, an identifier by which such metadata may be located, the location of the data unit **105** within a buffer, or any other suitable identifying information.

(36) The enqueue logic **120** resolves an incoming data unit **105** to one or more of the queues **140**

based on the destination port **190** of the data unit **105** and/or other suitable metadata. Enqueue logic **120** then attempts to enqueue the incoming data unit **105** within the resolved queue(s) **140**, meaning that the incoming data unit **105** is placed at the tail, or end, of the queue **140**. In some embodiments, enqueue logic **120** may be unable to enqueue a data unit **105** on account of the queue **140** being too large, having an unacceptable amount of delay, and/or other constraints, in which case the data unit **105** may be dropped instead of enqueued.

(37) In an embodiment, each queue **140** is assigned to a single port. The enqueue logic **120** will only enqueue a data unit **105** within a queue **140** assigned to a port **190** to which the data unit **105** is to be sent. As depicted, each port **190** is assigned to a single queue **140**. However, in other embodiments, there may be more than one queue **140** associated with some or all of the ports **190**. For instance, there may be different queues **140** assigned to a port **190** to store different types of data units **105**, to store data units **105** with different priority levels, and/or for other purposes.

(38) A data unit **105** may be enqueued in more than one queue **140** if it is to be transmitted to multiple ports **190**. For each port **190** to which a data unit **105** is to be sent, the enqueue logic **120** may utilize various rules to resolve a data unit **120** to a specific queue **140** associated with that port **190**.

(39) In an embodiment, the enqueue logic **120** is configured to link two or more small data units together when enqueueing the data units (e.g. using interpacket linking data). The enqueue logic **120** may do so, for instance, when rate adjustment logic has increased a maximum service rate for a port associated with the queue in which the data units are being enqueued, as described in other sections. As a consequence of such an action, in some embodiments, the linked data units may appear as a single data unit to dequeue logic **130** when it comes time to dequeue the linked data units, and thus be dequeued at the same time in manner that is transparent to the dequeue logic **130**.

DEQUEUE LOGIC

(40) Queue manager **110** further includes dequeue logic **130**. Dequeue logic **130** is configured to dequeue data units **105** from queues **140**. The process of dequeuing a data unit **105** from a queue **140** comprises various steps, depending on the embodiment, such as removing the data unit **105** from the head of the queue **140**, locating the data unit **105** in a buffer, and reading the dequeued data unit **105** from the buffer. A dequeued data unit **105** may further be removed from a corresponding buffer when if it is no longer in any of the queues **140**. The dequeue logic **130** further includes or is coupled to transmission logic for transmitting each dequeued data unit **105** to the appropriate port **190** or other downstream logic.

(41) Subsystem **100** is coupled to a clock (not depicted), whose signals indicate intervals of times referred to herein as clock cycles. Dequeue logic **130** may dequeue only a limited number of data units **105**, or even just one data unit **105**, per clock cycle. Depending on the embodiment, the dequeue process for a data unit may be completed in the same clock cycle in which it began, or the dequeue process may take place over a time period spanning multiple clock cycles, in which case it may be performed concurrently with other dequeue processes in a pipelined manner.

(42) The dequeue logic **130** is typically instructed (e.g. by the scheduler) as to the port(s) **190** for which a data unit **105** should be dequeued in a given clock cycle. The dequeue logic **130** identifies the queue(s) **140** associated with each indicated port **190**, and selects a specific one of those queues **140** from which to dequeue a data unit **105**. If more than one queue **140** is associated with a port **190**, the dequeue logic **130** may utilize any suitable selection mechanism to select between those queues **140** in a given clock cycle (e.g. round-robin, most filled, probability-based mechanisms, percentage-based mechanisms, etc.).

(43) 2.4. Scheduler

(44) Subsystem **100** further includes a scheduler **150**. Scheduler **150** instructs dequeue logic **130** when to dequeue a data unit for a given port **190** in accordance to a port scheduling algorithm. Selecting a port **190** for which a data unit should be dequeued in a given clock cycle is also referred to herein as servicing the port **190**. Although the examples described herein assume that the

scheduler **150** selects at most only one port for servicing per clock cycle, in other embodiments there may be two, three, or even greater numbers of ports **190** serviced in a given clock cycle.

(45) Any suitable scheduling algorithm may be utilized, including without limitation round-robin algorithms, algorithms based on random or pseudo-random functions, algorithms based on probabilities, percentages, or weights assigned to each port **190**, and so forth. In an embodiment, scheduler **150** intelligently services the ports **190** in such a manner that only ports **190** whose queues **140** have currently enqueued data are selected. That is, for instance, if queue **3** is empty (as depicted), the scheduler **150** will not select queue **3**. To this end, scheduler **150** receives queue control data **115** from queue manager **110**. Queue control data **115** may include a variety of information, depending on the embodiment, such as whether data is currently available in a queue **140** and/or for a port **190**, a count of the number of data units currently within a queue **140** and/or available for a port **190**, an indication that a data unit has been enqueued for a port **190** and/or queue **140**, an indication if a queue **140** assigned to a port **190** is ready to be dequeued, and/or other relevant queue data **115**. In an embodiment, queue control data **115** may further include one or more metrics, based on which the scheduler **150** may further adjust scheduling.

(46) In an embodiment, the scheduling algorithm is such that each port **190** for which data is enqueued will be selected an approximately equal number of times over a given time period. In an embodiment, the scheduling algorithm may be weighted, such that one port **190** is serviced more frequently than others. For instance, the number of times a given port **190** is selected over a period of time relative to other ports **190** with available data may be approximately proportional to the speed of the network interface associated with the given port **190** relative to the speeds of network interfaces associated with the other ports **190**. For instance, assuming port **0** is a 400 Gbps port, while port **1** is a 100 Gbps port, port **0** may be selected four times for each time port **1** was selected. In an embodiment, the scheduling algorithm may also or instead be weighted based on other factors, such as per-port performance metrics, queue sizes, or priority levels. These weights may change over time.

SEQUENCER

(47) According to an embodiment, scheduler **150** utilizes a sequencer **160** to help determine the order in which to schedule ports **190** for servicing. Sequencer **160** generates and/or makes use of a sequence **165**. The sequence **165** includes a plurality of slots, and each port **190** may be assigned one or more of the slots. For example, as depicted, the scheduler **150** utilizes a sequence **165** of six slots numbered, in order, 0 through 5.

(48) Although sequence **165** includes only one slot for each port **190**, a port **190** may be assigned more than one slot for various reasons. For example, certain ports **190** may be higher bandwidth ports **190**, and hence need to be scheduled more frequently than other ports **190**. For instance, supposing ports **0** and **1** are 400 Gb/s ports, while all other ports **190** are 100 Gb/s ports, ports **0** and **1** would need to be scheduled four times more frequently than the other ports **190**. The sequencer **160** could generate or utilize a variety of sequences to achieve this objective, such as, for instance, 0, 1, 2, 0, 1, 3, 0, 1, 4, 0, 1, 5.

(49) In an embodiment, the scheduler **150** selects the port **190** to service in an order based upon the sequence. For instance, the scheduler **150** may cycle through the slots in sequence **165** continually, selecting the next slot in the sequence **165** each clock cycle until the end is reached, at which point the sequence **165** begins anew. For example, using the depicted sequence **165**, the scheduler **150** may select port **0** in a first clock cycle, port **1** in a second, port **2** in a third, port **3** in a fourth, port **4** in a fifth, and port **5** in a sixth. The scheduler **150** may then return to port **0** for the seventh clock cycle, select port **1** in the eighth, and so on.

(50) Depending on the embodiment, the sequencer **160** may range from a relatively trivial component that returns a predefined sequence **160**, to a dynamic component that is constantly generating a sequence **165** out to some certain number of slots. In the former case, the sequence **165** may be fixed based on the hardware configuration of the ports **190**. However, in some

embodiments, port bandwidths may be at least partially configurable by a user, and hence the sequencer **165** may generate the sequence **165** programmatically based on the configuration (though the sequence **165** may not change much otherwise). In yet other embodiments, the sequencer **160** may intelligently generate the sequence **165** based on a variety of factors, such as device status data. The sequence **165** may be generated or regenerated at pre-defined intervals, or in response to various events or status-based triggers. The sequence **165** may be repeating or non-repeating, depending on the embodiment.

(51) In an embodiment, the sequencer **160** may at times rearrange the sequence **165** in an at least partially random manner, so as to introduce jitter into the scheduling process. A port **190** would still have the same number of slots in the sequence **165**, but the sequence **165** would change over time (e.g. 2,1,5,4,0,3 instead of 0,1,2,3,4,5).

PORT AVAILABILITY

(52) In an embodiment, the scheduler **150** selects the next slot in the sequence **165** whose corresponding port **190** is currently available for servicing. The availability of a port **190** may depend on several factors. For example, in an embodiment, a port **190** may be serviced only if data is actually enqueued for the port **190**. Whether or not data is enqueued for a port **190** may be specified directly in queue control data **115** received by the scheduler **150** each clock cycle, or scheduler **150** may keep track of enqueues and dequeues for the port **190** (as signaled in the queue control data **115**), so that it has an accurate count of the number of data units **105** in the port's queues **140**.

(53) For example, as depicted in FIG. 1, the queue **140** associated with port **3** is empty, and port **3** may thus be determined to be unavailable for servicing. The sequence **165** indicates that port **3** fills the slot between port **2** and port **4**—in other words, the scheduler **150** would normally select port **2** in a first clock cycle, port **3** in the second clock cycle, and port **4** in a third clock cycle. However, since port **3** has no data available, the second clock cycle would be wasted. Hence, the scheduler may skip port **3**, so that port **2** is selected in the first clock cycle, followed by port **4** in the second clock cycle.

SERVICE RATES

(54) The availability of a port **190** for service may further depend on other factors, such as service rate constraints. For example, hardware and/or logical constraints such as described elsewhere herein may make it undesirable or even impossible to service the same port **190** on back-to-back clock cycles, or more frequently than every third clock cycle, or at even lower rates, depending on the embodiment. A port **190** may thus be assigned a service rate that indicates how frequently the port **190** may be serviced. In some embodiments, all ports **190** have the same service rate. In other embodiments, there may be different service rates for different ports **190**.

(55) The scheduler **150** may include logic for enforcing a service rate. The scheduler **150** may store in its memory the identities of the port(s) **190** it has serviced in the most recent clock cycles. It may then decline to schedule a port **190** for the next clock cycle if doing so would result in the port **190** having a service rate higher than the service rate assigned to the port **190**. Note that it is possible that, in view of the service rate constraints and/or the lack of data in other queues **140**, no port **190** is serviced in certain clock cycles.

(56) For example, with the depicted sequence, suppose that only queues **2** and **3** have data currently enqueued therein. Further suppose that each port **190** has a maximum service rate of 1 in 3 clock cycles. The scheduler **150** would skip the first two slots in sequence **165** since the associated ports **0** and **1** have no data, and schedule port **2** in a first clock cycle, followed by port **3** in the second clock cycle. For the third clock cycle, the scheduler **150** would skip the last slots of the sequence **165**, since the associated ports **4** and **5** have no data, and again skip the first two slots of the sequence **165** for the same reason. Arriving at the slot for port **2**, the scheduler **150** would determine that port **2** cannot be serviced since it was already serviced within the last three clock cycles. For the same reason, the scheduler **150** would determine to skip port **3**, and hence no port

would be scheduled. For the fourth and fifth clock cycles, however, ports **2** and **3** could be scheduled, respectively, since doing so would no longer violate their respective service rates. Now, supposing in the interim, data was enqueued for port **5**. Port **5** would then be scheduled in the sixth clock cycle, and ports **2** and **3** could be scheduled for the seventh and eighth clock cycles without any service rate violations.

(57) As another example, suppose a sequence **165** of 0,1,0,2,0,3,0,4,5. Further suppose that, initially, only port **0** is occupied, and the each port has a maximum service rate of 1 in 2. Initially, the scheduler **150** would schedule port **0** every other clock cycle, and no port **190** would be serviced at other times. Supposing that ports **2** and **3** eventually enqueue data as well, the scheduler **150** would eventually fall into a pattern of 0,2,0,3,0,2,0,3 . . . , so long as ports **0**, **2**, and **3** have data.

(58) In an embodiment, rather than scheduler **150** having logic for enforcing service rates, scheduler **150** may instead receive a signal indicating when a port **190** is ready for new data, and transmission of this signal may be based on a service rate for the port **190**.

RATE ADJUSTER

(59) According to an embodiment, scheduler **150** includes rate adjustment logic, also referred to as rate adjuster **155**, configured to temporarily adjust the rate at which a port **190** is serviced responsive to certain events or conditions. For example, in some embodiments, such an event is a determination by dequeue logic **130** or other logic within queue manager **110** that certain smaller-sized data units are at the head(s) of certain queue(s) **140** associated with certain port(s) **190**, or otherwise ready for transmission to downstream logic. When such an event or condition is detected, and optionally while the condition continues to persist, the rate adjuster **155** may temporarily increase the service rate for the associated port **190** by two, or by some other factor, depending on the embodiment.

(60) The rate adjuster **155** may adjust the service rate for a port **190** in a variety of manners, depending on the embodiment. For instance, in some embodiments, the scheduler **150** does not actually change the service rate for calculation purposes, but rather additional mechanisms are enabled whereby the service rate realized by the scheduler **150** is effectively altered. Such mechanisms may include dequeuing multiple data units from a port simultaneously, temporarily storing one or more of those data units in a port FIFO, and/or enabling a pseudo port associated with that port FIFO within the sequence **165**. Examples of such mechanisms are described in other sections.

(61) In another embodiment, the rate adjuster **155** may simply lower the service rate used by the scheduler **150** when calculating when a port **190** may be re-serviced. Such an embodiment would rely on upstream and/or downstream logic to perform various measures to ensure that the port **190** could actually be serviced again within the time frame required. This may require the dequeue logic to enter into a special operation mode so that it may read data units and take appropriate follow-up measures more quickly. Or, this may require that other components transparently take certain measures to enable the faster service rate, such as preemptively dequeuing multiple data units at once or in response to a single request, linking such data units together prior to sending, separating such data units downstream, and so forth.

(62) The conditions or events that trigger a change in service rate for a port **190** may vary depending on the embodiment. In some embodiments, rate adjuster **155** increases the service rate for a port **190** at least partially in response to detecting that a small data unit is about to be dequeued for the port **190**. A small data unit may be, for instance, a data unit that is smaller than a threshold amount. The threshold may be fixed or user programmable. In an embodiment, the threshold is half the size of maximum size TDU. In an embodiment, a small data unit is any data unit that is smaller than the maximum TDU size. In an embodiment, in place of current data unit size, the rate adjuster may compare an average data unit size over a recent period of time to the threshold to determine when to adjust the service rate. In an embodiment, the rate adjuster may

adjust the service rate at least partially based on metadata from downstream logic (e.g. one or more metrics that indicate congestion at a particular downstream port or component, etc.).

(63) In some embodiments, rate adjuster **155** increases the service rate for a port **190** based at least in part on detecting that both the next data unit to be dequeued for the port **190** and the data unit that is second-in-line to be dequeued for the port **190** are small data units. The two data units may be considered small if they are both smaller than a threshold size, as described above, or if the two data units are, in aggregate, smaller than a threshold size. For instance, the rate adjuster **155** may increase the service rate for a port **190** if the next two data units are, in aggregate, smaller than the maximum TDU size (and thus can be transmitted together).

(64) In some embodiments, only EOP data units **105** can be considered small. The rate adjuster **155** will thus only trigger a rate adjustment for a port **190** if the next two or more data units are EOP data units (meaning that at least the second-in-line data unit is an SEOP data unit). In an embodiment, in addition to being EOP data units, the data units must still be, in aggregate, smaller than a certain size (e.g. the maximum TDU size). Or, in an embodiment, in addition to being EOP data units, the data units must still be smaller than a certain size individually (e.g. half the maximum TDU size).

(65) In an embodiment, the rate adjuster **155** may learn that the next data unit(s) for a port **190** are considered to be small based on queue control data **115** or other metadata from queue manager **110**, and thus adjust the service rate. The rate adjuster **155** may return to the default service rate upon receiving an indication that the next data unit(s) for the port **190** are no longer considered to be small.

(66) In other embodiments, when the next data unit(s) for a port **190** are considered to be small, the dequeue logic **130** is configured to automatically dequeue more than one data unit. The rate adjuster **155**, upon detecting that the dequeue logic **130** has dequeued multiple data units, takes other appropriate actions to adjust the service rate for the port **190**. The rate adjuster **155** returns to the original service rate when the dequeue logic **130** is no longer dequeue multiple data units for the port **190** (and/or when the multiple data units have been transmitted from to the port **190** from a port FIFO, as described in subsequent sections).

(67) In an embodiment, rate adjuster **155** may be separate from the scheduler **150**. For instance, the rate adjuster **150** may be implemented in queue manager **110** or in any other suitable component. The rate adjuster **155** may send instructions to the scheduler **150** to adjust the rate, and/or cause performance of any actions necessary to effect rate adjustment, in accordance to various techniques described herein.

(68) 2.5. Port FIFOS and Pseudo-Ports

(69) According to some embodiments, an increased service rate for a port **190** may be accomplished using port FIFOs, coupled with logic within the scheduler **150** that treats each port FIFO as a pseudo port. FIG. 2 illustrates a scheduler-based dequeue subsystem **200** configured to utilize port FIFOs **270**, according to an embodiment. Subsystem **200** may be utilized within any device that utilizes schedule-based dequeuing, including without limitation within the traffic manager or ingress arbiter of a network device. Subsystem **200** may be a part of subsystem **100**, or a separate subsystem. Conversely, subsystem **100** may, but need not necessarily, utilize the components and techniques described below with respect to subsystem **200**.

(70) Subsystem **200** includes dequeue logic **230** configured to dequeue queues **240** of data units, in a manner similar to that utilized by dequeue logic **130** in dequeuing queues **140**. Moreover, subsystem **200** includes a scheduler **250** configured to schedule ports **290** for servicing, in a manner similar to that utilized by scheduler **150** for scheduling ports **190** for servicing. Just as scheduler **150** utilizes a sequence **165** maintained by a sequencer **160**, scheduler **250** may utilize a sequence **265** maintained by a sequencer **260**.

DEQUEUEING MULTIPLE DATA UNITS TOGETHER

(71) In an embodiment, if the next two or more data units to be dequeued for a port **290** are both

small (e.g. each smaller than a threshold size, in aggregate smaller than a maximum TDU size, each EOP TDUs, etc.), dequeue logic **230** may dequeue more than one data unit for a port **290** at a time. Both data units are read and provided to the transmission logic as part of the same dequeue process. In such an embodiment, even though the rate at which the dequeue logic can perform dequeue events for a port **290** may not necessarily change, the rate at which data units are transmitted downstream is increased, because more than one data unit is dequeued at a time, thereby increasing the effective service rate for the port **290**.

(72) Depending on the embodiment, the dequeue logic **230** may dequeue multiple data units in different manners. For instance, rate adjustment logic may access control data or other metadata indicating the size and/or other characteristics of the next two data units in a queue. If the data units are both small, the rate adjustment logic may instruct the dequeue logic **230** to read and dequeue those data units in a single transaction.

(73) As another example, there may be two or more queues **240** assigned to a port. When the dequeue logic **230** dequeues the EOP data unit from one queue **240**, it may dequeue another data unit from another queue **240** assigned to the port **290** at the same time, if that data unit meets appropriate criteria (e.g. is a SEOP data unit smaller than a threshold size, small enough to be transmitted with the EOP data unit, etc.). In some embodiments, a port **290** may have a special queue **240** for SEOP data units, and the dequeue logic **230** may be configured to attempt to dequeue a second data unit from this queue whenever it dequeues an EOP data unit from another queue.

(74) In an embodiment, when dequeuing an EOP data unit, the dequeue logic **230** may proactively search for a second small data unit that is ready to be dequeued. For instance, the dequeue logic **230** may search for such a second data unit in the queue that was just dequeued, and/or the head data unit of each other queue assigned to the port. If such a second small data unit is found, the dequeue logic **230** may dequeue that data unit as well.

(75) In some embodiments, small data units may be linked together when enqueued (e.g. using interpacket linking data). Because the data units are linked, the dequeue logic **230** may not necessarily even know that it is dequeuing multiple data units.

PORT FIFOS

(76) In an embodiment, rather than automatically outputting all dequeued data units immediately downstream, the dequeue logic **230** includes or is coupled to special transmission logic **235**. The transmission logic **235** by default transmits dequeued data units to a port **290**, or to intermediate downstream logic, via a primary path **231**. However, the transmission logic **235** may in certain cases store dequeued data units in port FIFOs **270**, which are coupled to the ports **290** via secondary paths **232**.

(77) For example, in an embodiment, a port **290** and/or the downstream logic between dequeue logic **230** and port **290** may be capable of processing only one data unit for the port per clock cycle. Hence, the transmission logic **235** may transmit the first dequeued data unit out the primary path **231** to the appropriate port **290** or associated downstream logic **290**, and transmit the remaining dequeued data unit(s) out the secondary path **232** to the port FIFO **270** assigned to the port **290**. These data unit(s) are written to the port FIFO **270**, and dispatched to the port **290** in future clock cycles. Note that, since the data units buffered in a port FIFO **270** are eventually transmitted to the port **290** as well, dequeuing multiple data units may still effectively increase the service rate for the port **290** even when the port FIFO **270** is used.

(78) There may be a port FIFO **270** for each port **290**, though in some embodiments there may only be port FIFOs **270** for certain ports **290** (e.g. only high-speed ports **290**). Each port FIFO **270** comprises a number of entries. In an embodiment, each entry in the port FIFO **270** is of sufficient width to hold a full-size data unit. In an embodiment, each entry in the port FIFO **270** is of sufficient width to store a data unit whose size is no larger than whatever threshold may be utilized to classify a data unit as “small” for the purposes of determining when to increase the service rate

for a port **290**. The number of entries may vary depending on the embodiment, from two or even one, up to any desired maximum size.

(79) A port FIFO **270** may store a data unit in each entry. The port FIFO may use a first-in-first-out mechanism to organize the data units (e.g. a circular buffer), such that when reading from the port FIFO **270**, the oldest data unit will always be read first. When a data unit is read from a port FIFO **270**, the data unit is transmitted to the corresponding port **290**, or intermediate downstream logic, and the entry in which the data unit was stored is freed for storing another data unit

(80) A data unit in the port FIFO **270** for a port **290** is no longer considered to be enqueued within the queue(s) **140** that are associated with that port **290**. In an embodiment, the port FIFOs **270** are separate from any buffer in which data units are stored while they are enqueued in queues **140**.

(81) In an embodiment, if the port FIFO **270** for a port **290** already stores a data unit, the port FIFO **270** should be emptied before the transmission logic **235** transmits another data unit to the port **290** via the primary path **231**, so as to avoid reordering issues. The transmission logic **235** may thus transmit all data units dequeued for the port **290** to the port FIFO **270** until the port FIFO **270** is again empty.

(82) In an embodiment, dequeue logic **230** may only dequeue multiple data units for a port **290** if there is still room in the port FIFO **270** for that port **290**. Accordingly, dequeue logic **230** may receive a signal from port FIFO management logic indicating whether or not the dequeue logic **230** may dequeue multiple data units for the port **290** in the upcoming clock cycle. Of course, if there is no port FIFO **270** for a port **290**, the dequeue logic **230** should not dequeue more than one data unit for the port **290**.

(83) In some embodiments, such as when multiple dequeued data units are linked together or otherwise combined, the transmission logic **235** may include logic for splitting the dequeued data into multiple data units so that they may be transmitted along different paths **231/232**.

PSEUDO-PORTS

(84) Any suitable logic may be utilized to read and transmit data units from the port FIFOs **270**. However, in some embodiments, dispatch of data units from the port FIFOs **270** is also managed by the scheduler **250**. That is, the scheduler **250** determines when it is time for a port FIFO **270** to transmit a data unit stored therein, and instructs the port FIFO **270** to do so at the determined time. For any given clock cycle, the scheduler **250** may service a port **290** by instructing dequeue logic **230** to dequeue a data unit for that port **290** from a queue **240**, and/or instructing a port FIFO **270** for that port **290** to transmit a data unit.

(85) In an embodiment, the scheduler **250** may be configured to do so in a manner that effectively increases the service rate for the port **290** with minimal changes to the normal scheduling logic. The scheduler **250** may view each port FIFO **270** as a pseudo port. For instance, the port FIFO **270** for port **0** would be considered to be a pseudo port **0'**, the port FIFO **270** for port **1** would be considered a pseudo port **1'**, and so forth.

(86) The sequencer **260** may add these pseudo ports to the sequence **265**, so that each pseudo port has at least one slot in the sequence **265**, separate from the slot occupied by its corresponding port **290**. For instance, a sequence of ports **0-5** that includes pseudo ports might be **0,1,2,3,4,5,0',1',2',3',4',5'**. The scheduler **250** may select a pseudo port from the sequence **265** in the same manner it selects a normal port **290**. When a pseudo port is selected from the sequence **265**, instead of instructing the dequeue logic **230** to dequeue a data unit for the corresponding port **290**, the scheduler instructs the port FIFO **270** for the pseudo-port to transmit a data unit. In an embodiment, a pseudo port is only selected if the port FIFO **270** associated with that pseudo port is determined to be active (e.g. contains at least one data unit).

(87) In an embodiment, to avoid unfairly biasing a port **290** when its port FIFO **270** is active, the sequencer **260** may provide different sequences depending on whether a port FIFO **270** is currently active and/or whether there is currently data enqueued in the queue(s) **240** for that port **290**. The ratio between the number of slots assigned to each port **290** in the sequences **265**, including those

assigned to both the actual port **290** as well as its pseudo port, should remain the same regardless of whether a pseudo port is active. For instance, a default sequence **265** might be 0,1,2,3,4,5. When a port FIFO **270** is active for port **0**, the sequencer **265** might change the sequence to 0,1,2,3,4,5,0',1,2,3,4,5. Or, if there is not data in any queue **140** assigned to port **0**, the sequencer **265** might change the sequence to 0',1,2,3,4,5. The sequence **265** would revert to normal when the port FIFO **270** is no longer active.

(88) The underlying logic of the scheduler **250** need not actually change the service rate it enforces on a port **290**, but may instead enforce the service rate constraint on a pseudo port separately from its associated port **290**, or even not at all, in some embodiments. This results in increasing the effective service rate for the port **290**, transparently to the underlying scheduling algorithm. For example, if the maximum service rate for port **0** is every other clock cycle, the service rate would prevent the scheduler **250** from scheduling port **0** for back-to-back clock cycles, even if no data were enqueued for any other port. However, if port **0'** is active, the scheduler **250** could alternate between port **0** and port **0'** in back-to-back clock cycles, since neither port **0** nor port **0'** would exceed the maximum service rate. The effective service rate for port **0** would thus be doubled.

(89) As another example, suppose a sequence of ports **0,1,0,2,0,3,0,4**. With ports **0** and **2** available, the scheduler **250** would schedule ports according to the following pattern: 0,2,0,2,0,2, . . . and so forth. Now suppose that a number of small data units are dequeued consecutively for port **0**. As explained above, multiple of those data units may be dequeued concurrently, filling the port FIFO for port **0** (port **0'**). The sequence may thus change to 0,1,0',2,0,3,0',4 (or another suitable sequence). The scheduler **250** would then schedule ports according to the following pattern: 0,0',2,0,0',0,0',2,0,0' . . . and so forth.

(90) 2.6. Example Port FIFO Implementation for Single Port

(91) FIG. 3 is an illustrative view of components for a scheduler-based queueing subsystem **300**, with respect to an individual port that utilizes a port FIFO **370**, according to an embodiment. Subsystem **300** may, in some embodiments, be the same as subsystems **100** and/or **200**, but viewed from a different perspective. In other embodiments, subsystem **300** may be implemented without the specific details of FIG. 1 or 2, and subsystems **100** and **200** need not be implemented with the specific details of subsystem **300**.

(92) Subsystem **300** comprises enqueue logic **320** that receives requests to enqueue TDUs **305** destined for a port **0**. The enqueue logic **320** may be similar to enqueue logic **120**, for example. The requests are received at different times, and for each request, if permitted, the enqueue logic **120** adds the TDU identified in the request to a queue **340** that is associated with port **0**. There may be one or more queues **340** associated with port **0**, and any suitable selection mechanism may be utilized to select a specific queue **340** to add the TDU to. Queue(s) **340** may, in some embodiments, be similar to queues **140** and/or **240**.

(93) When enqueueing a TDU in a queue **340**, enqueue logic **320** may send a signal **325** to a scheduler **350** indicating that data is currently enqueued for port **0**. Based on the signal, scheduler **350** determines that port **0** is now active, in that there is data that needs to be dequeued for port **0**. In some embodiments, enqueue logic **320** may continue to send the signal so long as data remains enqueued for port **0**. In other embodiments, scheduler **350** may count the number of times it receives a signal **325**, and use that count to calculate how many TDUs are enqueued for port **0**.

(94) When port **0** is active, the scheduler utilizes scheduling logic, such as described with respect to scheduler **250** and/or elsewhere herein, to determine when it is time to dequeue the next TDU from queue(s) **340**. When the appropriate time arrives, the scheduler sends a signal **355** to dequeue logic **330** requesting that the dequeue logic **330** dequeue a TDU for port **0**. If the scheduler **350** maintains a count of the number of TDUs enqueued for port **0**, it may also decrement that count at this time.

(95) When the dequeue logic **330** receives the signal **355** from the scheduler **350**, the dequeue logic **330** selects one of the queues **340** associated with port **0**, if there are more than one queue **340** that currently enqueue data for port **0**. The queue **340** may be selected in any suitable manner. In an

embodiment, if any TDU other than a SOP TDU is at a head of one of the queues **340**, that queue **340** should be dequeued; otherwise, any queue **340** may be selected. Dequeue logic **330** may dequeue a TDU **332a** from the head of the selected queue **340** in any suitable manner, such as described with respect to dequeue logic **230** and/or **130**, or elsewhere herein.

(96) In some embodiments, dequeue logic **330** may, on some occasions, dequeue more than one TDU **332**. For example, in an embodiment, upon dequeuing a first TDU **332a** from the head of one queue **340**, if rate adjustment logic (e.g. in the dequeue logic **330** or at the scheduler **350**) determines that TDU **332a** meets certain first criteria, the rate adjustment logic may then check as to whether a next TDU **332b** may also be dequeued. The first criteria may be, for instance, that TDU **332a** is an EOP TDU and/or smaller than certain size (e.g. the maximum TDU size, a threshold TDU size, etc.).

(97) The rate adjustment logic may check for the next TDU **332b** at the head of the same queue **340** from which TDU **332a** was just dequeued, and/or from the head of a different queue **340**, depending on the embodiment. In order for dequeue logic **330** to dequeue the next TDU **332b**, the next TDU **332b** must meet second criteria. These second criteria may include one or more of: TDU **332b** is a SEOP TDU, TDU **332b** is smaller than a threshold size, and/or TDU **332b** and TDU **332a** combined are less than a threshold size (e.g. the maximum TDU size, the size of a certain bus, etc.). If the second criteria are met, TDU **332b** may be dequeued at the same time as TDU **332a**.

(98) In an embodiment, the rate adjustment logic may continue to look for additional TDUs **332** to dequeue with TDU **332a**, so long as the combined size of all dequeued TDUs is lower than a threshold size.

(99) In an embodiment, the dequeue logic **330** may dequeue multiple TDUs **332** in a manner that is transparent to the dequeue logic **330**. Rather, rate adjustment logic in the enqueue logic **320**, or in some other upstream logic, may have determined that TDUs **332a** and **332b** were both small TDUs, using logic similar to the above, and linked multiple TDUs together (e.g. at a SDU level), such that they appear to the dequeue logic **330** as a single TDU.

(100) As with dequeue logic **230**, dequeue logic **330** includes or is coupled to special transmission logic for outputting dequeued data units. This transmission logic includes a splitter **380** to which the dequeued data **332** is sent. Splitter **380** is configured to output a TDU **382** on a primary path **381** to downstream logic and/or send non-outputted TDU(s) **384** along a secondary path **383** to a port FIFO **370**, maintained by the transmission logic for temporary storage, depending on the state of subsystem **300**. In a default state, if the data output **332** by dequeue logic **330** comprises just one dequeued TDU **332a**, it passes through the splitter **380** as TDU **382**. A switch **385** controls the ultimate output of subsystem **300**. For any clock cycle in which switch **385** receives a TDU **382** on the primary path **381**, the switch **385** outputs the TDU **382**.

(101) Further in the default state, if the output **332** from dequeue logic **330** comprises two or more TDUs **332**, the splitter **380** splits a first TDU **382** from the remaining TDUs **384**. The first TDU **382** is again outputted via the primary path **381**, the one or more remaining TDUs **384** are stored as separate entries in the port FIFO **370**.

(102) However, if the port FIFO **370** is active, in that it stores one or more TDUs, the splitter **380** receives a signal **375** indicating that port FIFO **370** is active. It may continue to receive the signal **375** until the port FIFO **370** is inactive, or a separate signal may be sent when the port FIFO **370** is finally empty, depending on the embodiment. In any event, while the port FIFO **370** is active, the splitter does not output data via the primary path **381** at all. Rather, all of the data **332** is outputted along the secondary path **383** and stored in the port FIFO **370**.

(103) When the port FIFO **370** becomes active, it also sends signal **375** to the scheduler **350**, which interprets the signal **375** as indicating that data is now enqueued in the pseudo port **0'**. The scheduler **350** may thus determine whether port **0'** is active based on signal **375** in the same manner it determines whether port **0** is active from signal **325**. Whenever port **0'** is active, the scheduling logic of the scheduler **350** may schedule port **0'** for dequeuing in addition to or instead of port **0**. In

an embodiment, if the scheduling sequence does not already include port **0'**, the scheduler **350** may switch to a new scheduling sequence that includes port **0'**.

(104) When scheduler **350** schedules port **0'** for dequeuing, the scheduler sends a signal **356** to the port FIFO **370** indicating that the port FIFO **370** should release the TDU **372** at the head of the port FIFO. The port FIFO **370** outputs the TDU to the switch **385**. For any clock cycle in which the switch **385** receives a TDU **372** from the port FIFO **370**, the switch **385** outputs that TDU **372** from subsystem **300**. The signal **375** may also, in some embodiments, be sent to the switch **385**, so that it is simpler for the switch **385** to determine when to switch between the primary path **381** and the port FIFO **370**.

(105) In addition, the scheduler **350** may send signal **355** to port FIFO **370** at the same time it sends signal **355** to dequeue logic **330**. If the port FIFO **370** is active, the port FIFO **370** should output the TDU **372** at the head of the port FIFO **370**, since, to avoid ordering problems, any TDUs dequeued by dequeue logic **330** in that clock cycle will be added to the port FIFO **370** instead of transmitted along the primary path **381**.

(106) 2.7. Example Port FIFO Operation

(107) The operation of a port **0** and a pseudo-port **0'** for an associated port FIFO are now considered over a series of clock cycles. For simplification, it is assumed that all other ports are inactive during those clock cycles, and that port **0** is associated with a single queue. In each clock cycle, some or none of the following events may occur, depending on traffic patterns and other factors.

(108) First, a new data unit may be enqueued in the queue for port **0**.

(109) Second, a port may be scheduled for dequeuing, which is either port **0** (corresponding to the queue) or the pseudo-port **0'** (corresponding to the port FIFO). A service rate constraint prevents selection of port **0** in back-to-back clock cycles. However, in an embodiment, no service rate constraints apply to pseudo-port **0'**, and hence pseudo-port **0'** can be scheduled in back-to-back clock cycles. In other embodiments, there may be a service rate constraint on pseudo-port **0'**.

(110) In an embodiment, pseudo-port **0'** must be activated before it can be scheduled, and the port FIFO must store contents before pseudo-port **0'** can be activated. It may take a number of clock cycles after the port FIFO stores contents to actually activate the pseudo-port **0'** depending on the implementation. However, in other implementations, activation may occur immediately. Pseudo-port **0'** is deactivated when the port FIFO is empty.

(111) Third, if a port has been scheduled for dequeuing, a single data unit is outputted to port **0**. The outputted data unit may be forwarded directly to port **0**, or to intervening logic, such as a packet processor or traffic manager. If there are no contents in the port FIFO, and port **0** is scheduled, the data unit may be outputted directly from the head of the queue (the top of the queue contents column in the previous cycle). Otherwise, if port **0** or port **0'** is scheduled, a data unit is outputted from the head of the port FIFO (the top of the port FIFO contents column in the previous cycle).

(112) Fourth, if port **0** is scheduled, one or more data units may be moved from the head of the queue to the port FIFO. If a data unit is already being outputted directly from the queue to port **0**, only one data unit is moved to the port FIFO, and only if it is an SEOP data unit that can fit on the same bus as the other data unit. If the port FIFO already stores content, one or two data units may be moved from the head of the queue to the port FIFO when port **0** is scheduled. The number of data units moved to the port FIFO depends on their sizes and the number of data units in the queue (e.g. no bigger in aggregate than the maximum TDU size).

(113) In an embodiment, the depth of the port FIFO may be constrained. For instance, the port FIFO may be only two entries. The dequeue logic may be modified to avoid moving further data units to the port FIFO if the port FIFO has no room for the data units.

(114) 2.8. Miscellaneous

(115) Subsystems **100**, **200**, and **300** are merely examples of logic components for implementing

scheduler-based queueing within a network device or other computing device. Other implementations may include fewer and/or additional components in varying arrangements, and the distribution of work between components may likewise vary.

(116) In an embodiment, instead of being implemented in a system where the queues correspond to ports to which the data units are being sent, the queues may correspond to other constructs, such as ingress ports, traffic flows, traffic types, service classes, and so forth. The slots of sequence **160/260** may be assigned to these other constructs instead of ports.

3.0. Functional Overview

(117) This section describes example process flows for queueing and dequeuing data units in with an enhanced scheduler, such as described with respect to subsystems **100**, **200**, and/or **300** above. The flows may be performed within the switches and network devices described in other sections and/or in other computing devices. In an embodiment, each flow may be performed by a subsystem **100**, **200**, or **300**, though the flows may similarly be performed in other queueing subsystems, and subsystems **100**, **200**, and **300** need not be limited specifically to performance of the example flows described below.

(118) In an embodiment, each of the processes described in connection with the functional blocks described below may be implemented using one or more integrated circuits, logic components, computer programs, other software elements, and/or digital logic in any of a general-purpose computer or a special-purpose computer, while performing data retrieval, transformation, and storage operations that involve interacting with and transforming the physical state of memory of the computer.

(119) 3.1. Enqueueing a Data Unit

(120) FIG. **4** illustrates an example flow **400** for enqueueing a data unit in a queue, according to an embodiment. Block **410** comprises receiving a data unit, such as a TDU, PDU, or other data unit. In an embodiment, the data unit may be received at a traffic manager, ingress arbiter, or other component within a network device that needs to queue data units prior to outputting those data units to downstream logic. In some embodiments, in response to block **410**, the data unit may be placed in a buffer.

(121) Block **420** comprises determining a port associated with the data unit. The port may be determined in any suitable manner. For example, the port may be a destination port specified by control information passed along with the data unit from upstream logic (e.g. an ingress packet processor). Or, the port may be a destination port determined using forwarding logic based on address information and/or other suitable metadata within the data unit, or within another data unit belonging to a same PDU. As yet another example, the port may be an ingress port from which the data unit was received.

(122) Block **430** comprises identifying one or more queues assigned to the port. In some embodiments, there is only one queue assigned to the port, while in others there may be more. Different ones of these queues may, for instance, be assigned to hold data units having different characteristics, such as belonging to a specific traffic flow or group of flows, having a certain priority level, originating from a certain source, carrying a certain type of data, and so forth. There may also be multiple queues for load balancing or other purposes.

(123) Block **440** comprises selecting one of these queues to assign the data unit to, based on the characteristics of the data unit. Block **440** is, of course, a trivial step if there is only one queue for the port.

(124) Block **450** comprises determining whether the data unit may be enqueued in the selected queue. Under certain circumstances, the enqueue logic of a device may prevent a data unit from being enqueued. For instance, if the selected queue is above a certain queue size or exhibiting a certain amount of delay, the enqueue logic might determine not to enqueue the data unit.

(125) If the data unit may be enqueued, flow **400** proceeds to block **460**, which comprises enqueueing the data unit in the queue. The data unit is added to the end, or tail, of the queue,

typically by generating linking data that links the data unit (e.g. by reference to a location in which the data unit is stored) to the data unit that was previously at the tail of the queue, or generating new linking data for the queue if there is no data unit in the queue. For instance, in an embodiment, block **460** may be performed by enqueue logic, such as enqueue logic **120** or **320**, within the implementing component of the network device.

(126) If the data unit cannot be enqueued, flow **400** proceeds to block **470**, which comprises dropping the data unit without the data unit being enqueued in the selected queue. The data unit may be deleted from any buffer entry in which it was stored, if not enqueued elsewhere.

(127) Flow **400** is but one example method flow for enqueueing a data unit. Other flows may include fewer, additional, or different steps in varying arrangements. For instance, in an embodiment, multiple ports may be associated with the data unit, and blocks **420-440** are repeated for each port. Depending on the embodiment, various other steps may occur during the enqueue process, such as generating a signal to scheduling logic indicating that an enqueue event occurred, updating statistical counters or other information in the implementing device, and so forth.

(128) Flow **400** is repeated each time a data unit is received at the implementing component of the network device, resulting over a period of time in the enqueueing of many data units in multiple queues for multiple ports.

(129) 3.2. Scheduling Dequeue Operations

(130) A dequeue operation is used to dequeue a data unit that has been enqueued per flow **400**, so that it may be transmitted to the associated port and/or to downstream logic for further processing. Since enqueue operations may sometimes be performed more frequently than dequeue operations can be performed for a port, a scheduling process is utilized to determine when a given data unit from a given queue can be dequeued. Example dequeue process flows are described in other sections.

(131) FIG. **5** illustrates an example flow **500** for scheduling a dequeue operation, according to an embodiment. Flow **500** may be performed repeatedly (e.g. every clock cycle), so as to ensure that dequeue operations are performed as frequently as possible (e.g., again, each clock cycle). In some embodiments, flow **500** is performed one or more clock cycles in advance of the clock cycle in which the dequeue operation should begin, depending on the implementation. Flow **500** may be performed, for instance, by a scheduler **150**, **250**, or **350**, or other suitable scheduling logic.

(132) Block **510** comprises determining a scheduling sequence for the ports. The scheduling sequence is a sequence of slots, in which each slot is assigned to a different port, that indicates an order in which the ports may be selected. A port may have more than one slot in the sequence—for instance, the ratio of the number of slots assigned to a port relative to the number of slots assigned to another port may correspond to the ratio between their respective bandwidths. Sequences **165** and **265** are non-limiting examples of such a sequence.

(133) The sequence may be determined in any suitable manner. In an embodiment, the sequence should not usually change between iterations of flow **500**, but rather block **510** may be performed just once for any number of iterations of flow **500**. In an embodiment, the sequence may have been generated or selected by other logic asynchronously to the rest of flow **500**. Such logic may be triggered by certain conditions or events, such as activation of a port or pseudo-port, device configuration changes, changes in device state, etc. Or, the logic may repeatedly generate a schedule periodically or at other intervals based on the current device configuration and state. In such embodiments, block **510** may comprise accessing the sequence, or a portion thereof, in an appropriate memory location.

(134) Block **520** comprises determining which ports have data available for dequeuing. The ports may be determined in any suitable manner. In an embodiment, a scheduler may maintain port status information that indicates a count of the number of data units available for each port. This count may be updated in response to enqueue events and dequeue events. In another embodiment, queue management logic may send status information to the scheduler, indicating which queues and/or

ports currently have data enqueued. In yet another embodiment, a scheduler may have direct access to the queues, and be able to determine whether a port has data available by examining the queues directly.

(135) Block **530** comprises selecting the next port in the sequence. For instance, there may be a pointer that indicates the last sequence slot to have been selected (e.g. in previous iterations of flow **500**) or considered, or the first slot in the sequence if the sequence has recently been updated.

Block **530** may involve advancing the pointer and reading the port from the next assigned slot.

(136) Block **532** comprises determining if this port is available, based on the availability data. In other words, block **532** comprises determining whether any queue assigned to the port currently has data. If the port is available, then flow **500** proceeds to block **534**. Otherwise, flow **500** proceeds to block **536**.

(137) Block **534** comprises determining whether scheduling the port for clock cycle will exceed a service rate for the port. For instance, if the port has an assigned maximum service rate of one in every four clock cycles, this may involve determining whether the port has been scheduled in any of the previous three clock cycles. If the service rate will not be exceeded, flow **500** proceeds to block **540**. Otherwise, flow **500** proceeds to block **536**.

(138) If the port is available and does not exceed the assigned service rate, then, in block **540**, the port is scheduled for dequeuing. This may involve, for instance, sending an instruction to dequeue a data unit in a queue associated with the port to dequeue logic for the upcoming clock cycle.

(139) Flow **500** then proceeds to block **550**, which comprises dequeuing a data unit from a queue associated with the port, as scheduled. The queue may be any queue associated with the port, if there are more than one queues associated with the port. The queue may be selected by the dequeue logic using any suitable means, or the scheduling logic may specify the queue to release a data unit from. In an embodiment, the selected queue is the queue from which a data unit was most recently released for the port, if the previously released data unit was not an EOP TDU. Otherwise, any suitable selection algorithm may be used to select the queue, including round-robin, random, priority-based, and so forth.

(140) Dequeuing the data unit may involve a number of steps, depending on the embodiment, such as “popping” a node representative of the data unit from the head of the queue by manipulating linking data representative of the queue, locating one or more buffer addresses that store the data unit based on address information within the node, reading the data unit from those address(es), freeing the address(es) for other data, and so forth.

(141) Flow **500** the proceeds to block **560**, which comprises transmitting the data unit to downstream logic, such as a packet processor, traffic manager, transmit buffer, port, etc., depending on the component implementing flow **500**. Note that, in an embodiment, blocks **540-560** may be referred to, collectively, as servicing the associated port. The process of servicing the port may actually span multiple clock cycles, and the clock cycle in which it was “scheduled” may simply be the clock cycle in which the process of dequeuing the data unit begins. Note that the dequeuing of data units may be pipelined, so that multiple data units are in various stages of being dequeued concurrently.

(142) On the other hand, if the port is not available, or if the service rate for the port will be exceeded, the port cannot be serviced in the upcoming clock cycle. Flow **500** may thus proceed to block **536**, which comprises determining whether there are any more ports in the sequence. If so, flow **500** returns back to block **530** for consideration of the next port in the sequence. Otherwise, flow **500** proceeds to block **570**, in which the scheduling process terminates without a port having been scheduled for servicing in the upcoming clock cycle.

(143) Flow **500** is but one example method flow for scheduling dequeue operations. Other flows may include fewer, additional, or different steps in varying arrangements. For instance, in an embodiment, rather than determining port availability for all ports up front in block **520**, port availability is only determined when necessary in block **532**. Moreover, blocks **532** and **534** may be

reversed in order. In an embodiment, rather than following a fixed sequence, a ranking or other mechanism may be utilized to select the next port in block **530** based on other logic and/or factors. Depending on the embodiment, various other steps may occur during the dequeue process, such as generating a signal to scheduling logic indicating that a dequeue event occurred, updating statistical counters or other information in the implementing device, and so forth.

(144) Moreover, as described subsequently, in some embodiments, multiple smaller data units may be dequeued simultaneously in block **550**, though their transmission may in some cases be staggered through the use of a FIFO, cache, or other mechanism.

(145) 3.3. Increasing Service Rate for Small Data Units

(146) FIG. **6** illustrates an example flow **600** for increasing the service rate of a port for smaller data units, according to an embodiment. Flow **600** may be performed, for instance, by dequeuing logic and/or scheduling logic within a network device, such as described in other sections.

(147) Block **610** comprises determining characteristics of upcoming data units to be dequeued for a port. The data units may be, for instance, one or two data units at the head of the next queue from which the dequeue logic will dequeue a data unit for the queue. For instance, if a PDU is currently in the process of being dequeued, the next queue may be the queue that contains the next data unit for that PDU. In an embodiment, the queues may include any number of queues, or even all queues, assigned to the port.

(148) The characteristics may be determined using any suitable data. For instance, block **610** may comprise reading control data or other metadata for the data unit. In an embodiment, the determined characteristics may include the size of each of the data units. For instance, block **610** may entail determining that the next data unit to be dequeued is 64 bytes, and the data unit following it is 128 bytes. Or block **610** may involve determining that at least two different queues each have 64 byte data units at their heads. In an embodiment, rather than determining an actual size of the data units, it is only determined whether or not the data units are full-size data units, as described in other sections. In an embodiment, block **610** may also or instead comprise classifying a data unit as an SOP, EOP, MOP, and/or SEOP data unit, based on its position in a PDU.

(149) Block **620** comprises determining whether the next data units are smaller than a threshold size. For instance, the threshold may be half the size of a full-sized TDU, and block **620** may comprise comparing the sizes of each of the next two data units to be dequeued to that threshold individually. In an embodiment, instead of being based on the sizes of the next data units individually, the determination may be based on the sizes of the next data units to be dequeued, combined. For instance, the threshold may be the size of a full-sized TDU, and the cumulative size of the next two data units to be dequeued may be compared to this size.

(150) In an embodiment, blocks **610** and **620** may be performed as part of process of attempting to identify a set of data units to dequeue next, such that the set is guaranteed to be no greater than a full-sized TDU. Such a process may be triggered, for instance, upon determining that the next data unit to be dequeued is the EOP of a PDU that is currently in the process of being dequeued, or that the next data unit to be dequeued is a SEOP. Upon such a determination, scheduling logic or dequeuing logic may examine one or more other queues (or the next data unit in the same queue), to determine if any data unit at the head of one of those queues is smaller than a certain size (e.g. the difference between the next data unit to be dequeued and a full-size TDU).

(151) If the next data units are smaller than the threshold, then in block **630**, the port is served at increased maximum service rate. The service rate may be, for instance, doubled, increased by 50%, or any other suitable adjustment. In an embodiment, the service rate may be increased directly, in that a maximum service rate configuration setting that controls the maximum frequency with which scheduling logic selects a port for servicing is explicitly increased (e.g. increasing the service rate used in block **540** from every other clock cycle to, for instance, every clock cycle). In another embodiment, the service rate may be increased indirectly in a manner that is transparent to the underlying scheduling logic. This may be accomplished by, for instance, causing the dequeuing

logic to dequeue multiple data units at once, activating a pseudo-port within a scheduling sequence, activating a port FIFO, and/or taking other measures as described elsewhere herein.

(152) If, on the other hand, the next data units are larger than the threshold, then in block **640**, the port is served at its default maximum service rate. In some embodiments, if the maximum service rate for the port is at an increased value due to previous iterations of flow **610** having branched to block **630**, then block **610** may involve returning to the default service rate. For example, in an embodiment, this may involve returning a maximum service rate configuration setting to its default value. Or, this may involve disabling, or not activating, various features used to increase the maximum service rate indirectly. For instance, in embodiments where multiple data units may be dequeued to increase the service rate, block **640** may involve dequeuing only one data unit instead of multiple data units.

(153) Flow **600** is but one example method flow for increasing the service rate of a port. Other flows may include fewer, additional, or different steps in varying arrangements. For example, although examples are given herein where only two data units are considered for dequeuing in block **630**, in other embodiments, more than two data units may be dequeued if there is sufficient bandwidth (e.g. if the next three data units to be dequeued are cumulatively smaller than the maximum TDU size). Moreover, flow **600** may be repeated any number of times. In an embodiment, flow **600** may be performed in advance of each clock cycle.

(154) In an embodiment, rather than being based on the size of the data unit alone, the determination of block **620** may be but one factor in a more complex algorithms for determining whether to branch to block **630** or **640**. For instance, the algorithm may include a comparison of the average data unit size over a recent period of time to the threshold (or a different threshold, or a determination of whether a queue size or average queue size exceeds some threshold. Additional considerations for such an algorithm may be based on any other suitable metadata. For instance, such metadata may be received from downstream logic, such as the fill level of a buffer associated with a downstream component, an effective transmission rate of a port, a congestion flag, or any other suitable data. In an embodiment, such an algorithm need not necessarily include a comparison of the actual data unit size to a threshold.

(155) 3.4. Dequeuing a Normal Port Coupled to a Port FIFO

(156) FIG. 7 illustrates an example flow **700** for implementing a dequeue operation in response to a request to dequeue a normal port that is coupled to a port FIFO, according to an embodiment. Flow **700** may be performed, for example, with respect to subsystems **200** and/or **300**, as well as other queueing subsystems that utilize a port FIFO.

(157) Block **710** comprises receiving an instruction to dequeue a data unit for a port. Such an instruction may be received, for instance, from scheduling logic that selects the port for servicing, such as, without limitation, in performance of block **540** of FIG. 5.

(158) Block **720** comprises determining whether the port FIFO for the specified port is empty. If so, flow **700** proceeds to block **725**, which comprises outputting the first data unit in one of the port's queues to downstream logic, such as to a packet processor, traffic manager, transmit buffer, etc. This may include various substeps, such as described with respect to blocks **550-560** of FIG. 5.

(159) If, on the other hand, the port FIFO is not empty, then the data to output downstream should come from the port FIFO rather than the queue, so as to avoid ordering issues. Flow **700** thus proceeds to block **730**. Block **730** comprises outputting the first data unit in the queue to the port FIFO instead of the downstream logic. Block **735** then comprises outputting the data unit at the head of the port FIFO to the downstream logic. Blocks **730** and **735** may be performed concurrently, or in any order.

(160) From blocks **725** and **735**, flow **700** proceeds to block **740**. Block **740** comprises determining whether the next data unit in the queue may also be dequeued for the same time slot (e.g. the same clock cycle, the same set of clock cycles used to dequeue the first data unit, etc.). In an embodiment, the next data unit may also be dequeued if both it and the first data unit (that was

outputted in block **725** or **730**) are smaller than a certain size, individually, or in combination. In an embodiment, the next data unit may be dequeued if it is an SEOP and the first data unit was an SOP. In an embodiment, the next data unit may be dequeued if it is smaller than the bandwidth remaining on a transmission bus from the queue after the first data unit is outputted.

(161) If the next data unit can also be dequeued, then flow **700** proceeds to block **750**, which comprises determining whether the port FIFO is full. The port FIFO is considered full if it does not have room to store the next data unit, and/or has a predefined maximum number of data units already stored within it, depending on the embodiment. If the port FIFO is not full, then the next data unit is outputted to the port FIFO in block **760**.

(162) Then, in block **765**, the port's pseudo-port is activated, if it is not already active. This may comprise, for instance, adding the pseudo-port to a list of active ports considered by the scheduling logic, selecting an updated scheduling sequence that includes the pseudo-port, or taking any other steps described herein.

(163) If, on the other hand, it was determined in block **740** that the next data unit cannot also be dequeued, or if the port FIFO is full in block **750**, then the next data unit is not outputted from the queue at this time, and flow **700** proceeds to block **770**. Block **770** comprises determining whether the port FIFO is empty. If so, then the port's pseudo-port is deactivated.

(164) Flow **700** is but one example method flow for dequeuing a port. Other flows may include fewer, additional, or different steps in varying arrangements. For example, blocks **740** and **750** may be reversed. As another example, the various determinations of flow **700** may actually be performed in advance of the time slot in which the data unit(s) are outputted, based on the expected states of the associated queue and port FIFO. The data unit(s) may then be outputted concurrently in the same clock cycle or other slot of time, in accordance to the determinations previously made.

(165) As another example, a port may have more than one queue assigned to it. Flow **700** may thus include a step of determining the queue from which to dequeue the first data unit. In an embodiment, the queue of the next data unit is the same as this queue. In another embodiment, the next data unit may be selected from a different queue, and hence flow **700** may comprise selecting the different queue. In an embodiment, the next queue may be selected based on the sizes of the data units at the heads of the queues assigned to the data unit, so as to increase the likelihood that the next data unit can be dequeued.

(166) In yet another embodiment, flow **700** may loop back to **740** from block **760** or **765** for consideration of whether a third, fourth, or other additional data unit may also be outputted to the port FIFO in the same time slot.

(167) 3.5. Dequeuing a Pseudo Port for a Port FIFO

(168) According to an embodiment, activation of a port's pseudo-port, such as occurs in block **765**, indirectly enables servicing of the port at an increased effective maximum service rate, since both the port and the pseudo-port may be scheduled independently, but the data outputted from servicing the port and the pseudo-port originates from the same set of queues. Conversely, once the pseudo-port is deactivated, the effective maximum service rate for the port returns to its default level by virtue of the pseudo-port no longer being available for scheduling.

(169) FIG. **8** illustrates an example flow **800** for implementing a dequeue operation in response to a request to dequeue a pseudo-port associated with a port FIFO, according to an embodiment. In some embodiments, flow **800** may be utilized in conjunction with flows **500** and/or **700**. For example, flow **700** may be performed when the scheduling logic of flow **500** selects the port for servicing, while flow **800** may be performed when the port's pseudo-port is scheduled for servicing. However, flows **500** and **700** may also be performed without the specific details of flow **800**.

(170) Block **810** comprises receiving an instruction to dequeue a data unit for the pseudo-port. Such an instruction may be received from, for instance, from scheduling logic that selects the port for servicing, such as, without limitation, in performance of block **540** of FIG. **5**. Receipt of such an instruction may only be possible when the pseudo-port is activated, such as may occur as a

result of block **765** in flow **700**, or in response to other events.

(171) Block **820** comprises, in response to the instruction, outputting the data unit at the head of the port FIFO to downstream logic, such as described with respect to block **735** of flow **700**.

(172) Block **830** comprises determining whether the port FIFO is empty. If so, then in block **840**, the pseudo-port is deactivated, as in block **775** of flow **700**.

(173) Flow **800** is but one example method flow for dequeuing a pseudo-port. Other flows may include fewer, additional, or different steps in varying arrangements.

4.0. Example Embodiments

(174) Examples of some embodiments are represented, without limitation, as follows:

(175) According to an embodiment, a network apparatus comprises: a plurality of ports configured to send and/or receive data units over one or more networks; enqueue logic configured to enqueue the data units in queues, in which the data units await transmission to one or more downstream logic components; dequeue logic configured to dequeue particular data units from the queues in response to instructions from scheduling logic; transmission logic configured to send the particular data units that have been dequeued by the dequeue logic to the one or more downstream logic components; the scheduling logic, configured to determine, at a given time, a given port of the ports to service, and send an instruction to the dequeue logic to dequeue data associated with the given port at the given time, the scheduling logic configured to service at least a first port of the ports no more frequently than at a defined service rate; rate adjustment logic configured to, based on sizes of certain data units, of the data units, that are assigned to the first port, and/or based on metadata received from the one or more downstream logic components, temporarily adjust the service rate from a first rate to a second rate.

(176) In an embodiment, the rate adjustment logic is configured to adjust the service rate from the first rate to the second rate responsive to a determination that at least two data units associated with the first port are both ready to be dequeued and smaller than a particular size.

(177) In an embodiment, determining that at least two data units associated with the first port are smaller than the particular size comprises at least one of: determining that the at least two data units are, in combination, smaller than a maximum data unit size.

(178) In an embodiment, determining that at least two data units associated with the first port are smaller than the particular size comprises at least one of: determining that the at least two data units are, in combination, smaller than a user-programmable size.

(179) In an embodiment, the rate adjustment logic is configured to adjust the service rate from the first rate to the second rate responsive to a determination that an average size of a plurality of the data units over a recent period of time is smaller than a particular size.

(180) In an embodiment, the rate adjustment logic is configured to adjust the service rate from the first rate to the second rate responsive based on information conveyed to the rate adjustment logic from the one or more downstream logic components.

(181) In an embodiment, the first rate is a rate at which the first port must be serviced to ensure that a defined bandwidth of the first port is realized when the first port is transmitting maximum-size data units, wherein the second rate is greater than the first rate.

(182) In an embodiment, the first rate is once per every other clock cycle, wherein the second rate is once per clock cycle, wherein the bandwidth is at least 800 Gbps.

(183) In an embodiment, the rate adjustment logic is configured to adjust the service rate by at least causing the dequeue logic to dequeue at least two data units for the first port concurrently at a particular time.

(184) In an embodiment, the rate adjustment logic is configured to adjust the service rate by at least causing the enqueue logic to link two or more data units together when enqueueing the two or more data units in a queue associated with the first port;

(185) In an embodiment, the dequeue logic to dequeue the two or more data units from the queue concurrently at a particular time in a manner that is transparent to the dequeue logic as a

consequence of the linking.

(186) In an embodiment, the rate adjustment logic is configured to adjust the service rate by at least causing the dequeue logic to dequeue multiple data units for the first port at a same time, wherein the transmission logic is configured to only send one data unit for the first port to the one or more downstream logic components at a time, wherein the transmission logic includes a port FIFO for the first port, the port FIFO configured to store at least one of the multiple data units for transmission to the one or more downstream logic components at a later time.

(187) In an embodiment, when the port FIFO stores at least one data unit, the rate adjustment logic is configured to activate a pseudo-port for the first port in a port scheduling sequence for the scheduler, the first port and the pseudo-port separately selectable from the sequence at a frequency corresponding to the first rate, the first port thereby effectively serviced at the second rate.

(188) In an embodiment, when the port FIFO stores at least one data unit, and the scheduler determines to service the first port or a corresponding pseudo-port in a given clock cycle, the scheduler is configured to instruct the transmission logic to transmit a next data unit in the port FIFO to the one or more downstream logic components instead of a newly dequeued data unit.

(189) In an embodiment, each queue of the queues is assigned to exactly one of the ports, the queue holding only data units that are to be sent to or were received from the port assigned to the queue, the dequeue logic configured to, when instructed to dequeue a data unit associated with a particular port, select a particular queue of the queues assigned to the particular port from which to dequeue a particular data unit at the head of the particular queue.

(190) In an embodiment, when the particular data unit is an end-of-packet data unit and the rate adjustment logic detects another end-of-packet data unit at the head of another queue assigned to the particular port, or immediately following the particular data unit in the particular queue, the dequeue logic is configured to dequeue both the particular data unit and the other end-of-packet data unit at the same time.

(191) In an embodiment, the apparatus further comprises the one or more downstream logic components, the one or more downstream logic components being one or more traffic managers, one or more packet processors, or the plurality of ports.

(192) In an embodiment, the queues are ingress queues and the one or more downstream logic components is a traffic manager. In an embodiment, the queues are egress queues maintained by a traffic manager that is buffering the data units in a buffer memory, and the one or more downstream logic components are packet processors. In an embodiment, a data unit is ready to be dequeued when it is at a head of a queue assigned to a port associated with the data unit, or when the data unit immediately follows an end-of-packet data unit at the head of the queue.

(193) According to an embodiment, a method comprises: sending and/or receiving data units over one or more networks via a plurality of ports of a network apparatus; enqueueing the data units in queues, in which the data units await transmission to one or more downstream logic components of the network apparatus; dequeuing particular data units from the queues, the dequeuing including determining, at a given time, a given port of the ports to service, based at least on a constraint that a first port of the ports is serviced no more frequently than at a defined service rate, and dequeuing data associated with the given port at the given time; sending the particular data units that have been dequeued to the one or more downstream logic components; based on sizes of certain data units, of the data units, that are assigned to the first port, temporarily adjusting the service rate from a first rate to a second rate.

(194) In an embodiment, adjusting the service rate from the first rate to the second rate is responsive to a determination that at least two data units associated with the first port are both ready to be dequeued and smaller than a particular size or that an average size of a plurality of the data units over a recent period of time is smaller than a particular size.

(195) In an embodiment, the first rate is a rate at which the first port must be serviced to ensure that a defined bandwidth of the first port is realized when the first port is transmitting maximum-size

data units, wherein the second rate is greater than the first rate.

(196) In an embodiment, adjusting the service rate comprises dequeuing the at least two data units concurrently at a particular time.

(197) In an embodiment, adjusting the service rate comprises dequeuing multiple data units for the first port at a same time, wherein only one data unit is sent for the first port to the one or more downstream logic components at a time, wherein the method further comprises: storing at least one of the multiple data units in a port FIFO for the first port; when the port FIFO stores at least one data unit, activating a pseudo-port for the first port in a port scheduling sequence; wherein determining the given port to service at the given time comprises selecting the given port using the port scheduling sequence, the first port and the pseudo-port separately selectable from the sequence at a frequency corresponding to the first rate, the first port thereby effectively serviced at the second rate; and when the port FIFO stores at least one data unit, upon determining to service the first port or a corresponding pseudo-port in a given clock cycle, transmitting a next data unit from the port FIFO to the one or more downstream logic components, instead of a newly dequeued data unit.

(198) In an embodiment, each queue of the data unit queues is assigned to exactly one of the ports, the queue holding only data units that are to be sent to or were received from the port assigned to the queue; wherein the method further comprises, when dequeuing data associated with a particular port, selecting a particular queue of the queues assigned to the particular port from which to dequeue a particular data unit at the head of the particular queue; when the particular data unit is an end-of-packet data unit and there is another end-of-packet data unit at the head of another queue assigned to the particular port, or that immediately follows the particular data unit in the particular queue, dequeuing both the particular data unit and the other end-of-packet data unit at the same time.

(199) In an embodiment, the one or more downstream logic components is one or more traffic managers, one or more packet processors, or the plurality of ports.

(200) According to an embodiment, one or more non-transitory computer-readable media store instructions that, when executed by one or more computing devices, cause performance of the foregoing.

(201) Yet other example embodiments are described in other sections herein.

5.0. Example Network Device

(202) This section describes, in greater detail, example network devices in which the components and techniques described herein may be implemented. However, the techniques described herein are also useful in switches and contexts other than those described in this section.

(203) 5.1. Networks

(204) FIG. **9** is an illustrative view of various aspects of an example networking system **900**, also referred to as a network, in which the techniques described herein may be practiced, according to an embodiment. Networking system **900** comprises a plurality of interconnected nodes **910a-910n** (collectively nodes **910**), each implemented by a different computing device. For example, a node **910** may be a single networking computing device, such as a router or switch, in which some or all of the processing components described herein are implemented in application-specific integrated circuits (ASICs), field programmable gate arrays (FPGAs), or other integrated circuit(s). As another example, a node **910** may include one or more memories storing instructions for implementing various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by the various components.

(205) Each node **910** is connected to one or more other nodes **910** in network **900** by one or more communication links, depicted as lines between nodes **910**. The communication links may be any suitable wired cabling or wireless links. Note that system **900** illustrates only one of many possible

arrangements of nodes **910** within a network. Other networks may include fewer or additional nodes **910** having any number of links between them.

(206) While each node **910** may or may not have a variety of other functions, in an embodiment, each node **910** is configured to send, receive, and/or relay data to one or more other nodes **910** via these links. In general, data is communicated as series of discrete units or structures of data represented by signals transmitted over the communication links.

(207) When a node **910** receives a unit of data, it typically examines addressing information within the unit of data (and/or other information within the unit of data) to determine how to process the unit. The addressing information may be, for instance, a Media Access Control (“MAC”) address, Internet Protocol (IP) address, MPLS label, and/or any other suitable information. If the addressing information indicates that the receiving node **910** is not the destination for the data unit, the receiving node **910** may look up the destination node **910** within the receiving node's routing information and route the data unit to another node **910** connected to the receiving node **910** based on forwarding instructions associated with the destination node **910** (or an address group to which the destination node belongs). The forwarding instructions may indicate, for instance, an outgoing port over which to send the unit of data, a label to attach the unit of data, etc. In cases where multiple paths to the destination node **910** are possible, the forwarding instructions may include information indicating a suitable approach for selecting one of those paths, or a path deemed to be the best path may already be defined.

(208) Addressing information, flags, labels, and other metadata used for determining how to handle a data unit are typically embedded within a portion of the data unit known as the header. The header is typically at the beginning of the data unit, and is followed by the payload of the data unit, which is the information actually being sent in the data unit. A header is typically comprised of fields of different types, such as a destination address field, source address field, destination port field, source port field, and so forth. In some protocols, the number and the arrangement of fields may be fixed. Other protocols allow for arbitrary numbers of fields, with some or all of the fields being preceded by type information that explains to a node the meaning of the field.

(209) Different nodes **910** within a network **900** may send, receive, and/or relay data units at different communication levels, or layers. For instance, a first node **910** may send a unit of data at the transport layer (e.g. a TCP segment) to a second node **910** over a path that includes an intermediate node **910**. This unit of data, which may sometimes first be broken into smaller units, may be encapsulated in units of data at various sublevels before it is transmitted from the first node **910**. For example, a TCP segment may be encapsulated into IP packets, those IP packets may be encapsulated into Ethernet frames, and those Ethernet frames may eventually be sent out as a collection of signal-encoded bits to the intermediate device. Depending on the network layer at which the device operates, the intermediate node **910** may examine or, and in some cases reconstruct, the original frames, packets, and/or segment before routing that data to the second node **910**. Or, the intermediate node **910** may simply examine the frames, or in some cases the packets as well, and route the data to the second node **910** without ever examining the original segment (or, in some cases, the packet).

(210) Data units of the network layer at which a node **910** operates are considered protocol data units (“PDUs”), such as, without limitation, IP packets for the network layer (layer 3), frames for the Ethernet layer (layer 2), etc. A PDU may, for data unit handling purposes both within and without the device, be broken up into smaller subunits, or portions. The node **910** need not necessarily ever assemble the data in a PDU together, but rather may in an embodiment act upon constituent portions of the PDU, which may be linked together by identifiers, linked lists, or other suitable constructs. Within a node **910** itself, these portions are referred to herein as transport data units (“TDUs”). For instance, if the PDU is a packet, the TDU might be one or more cells. The first TDU in a PDU is referred to as the start-of-packet (“SOP”), while the last TDU in the PDU is referred to as the end-of-packet (“EOP”).

(211) Generally speaking, the TDU is the largest contiguous unit of data that certain internal components of a node **910** are configured to communicate between each other in a given period of time. For instance, a node **910** may have a traffic manager that is capable of receiving no more than a single TDU from each input interface it is connected to during a single clock cycle of the traffic manager. Additionally, though not necessarily the case in all embodiments, in at least some embodiments, the contiguous portions of data sent by each port of a node **910** in a given period of time may be no larger than a TDU. In an embodiment, each TDU is of a fixed size, except for the last TDU in a PDU, which may be of a size less than the fixed size.

(212) In some embodiments, for physical storage purposes, a TDU may further be divided into chunks referred to as storage data units (“SDUs”). In an embodiment, an SDU is the largest contiguous portion of data that may be stored in a physical buffer entry of a particular buffer. In other words, the maximum size of an SDU is the same as the maximum size of a physical buffer entry. In an embodiment, the maximum number of SDUs in a TDU is fixed. However, an EOP TDU may have less than this number of SDUs. Moreover, the last SDU in a TDU (e.g. the EOP TDU) may be smaller than maximum SDU size.

(213) In an embodiment, TDU and SDU boundaries may be relative to the component acting upon the data. That is, for example, a node **910** whose traffic manager is configured to use TDUs of a first size and SDUs of a second size may further include other components configured to communicate or buffer data units of sizes other than the first size and the second size.

(214) For convenience, many of the techniques described in this disclosure are described with respect to embodiments where the PDUs are IP packets in a L3 (level 3) network, and the TDUs are constituent cells and/or frames in an L2 (level 2) network, in which contexts the described techniques have particular advantages. It will be recognized, however, that these techniques may also be applied to realize advantages in routing other types of data units conforming to other protocols and/or at other communication layers within a network. Thus, unless otherwise stated or apparent, the techniques described herein should also be understood to apply to contexts in which the PDUs, TDUs, and SDUs are of any other types of data structures, such as segments, InfiniBand Messages, or datagrams. That is, in these contexts, other types of data structures may be used in place of packets, cells, frames, and so forth.

(215) Any node in the depicted network **900** may communicate with any other node in the network **900** by sending data units through a series of nodes **910** and links, referred to as a path. For example, Node B (**910b**) may send data units to Node H (**910h**) via a path from Node B to Node D to Node E to Node H. There may be a large number of valid paths between two nodes. For example, another path from Node B to Node H is from Node B to Node D to Node G to Node H.

(216) In an embodiment, a node **910** does not actually need to specify a full path for a data unit that it sends. Rather, the node **910** may simply be configured to calculate the best path for the data unit out of the device (e.g. which egress port it should send the data unit out on). When a node **910** receives a data unit that is not addressed directly to the node **910**, based on header information associated with a data unit, such as path and/or destination information, the node **910** relays the data unit along to either the destination node **910**, or a “next hop” node **910** that the node **910** calculates is in a better position to relay the data unit to the destination node **910**. In this manner, the actual path of a data unit is product of each node **910** along the path making routing decisions about how best to move the data unit along to the destination node **910** identified by the data unit.

(217) 5.2. Network Device

(218) FIG. **10** is an illustrative view of various aspects of an example network device **1000** in which techniques described herein may be practiced, according to an embodiment. Network device **1000** is a computing device comprising any combination of hardware and software configured to implement the various logical components described herein, including components **1010-1090**. For example, the apparatus may be a single networking computing device, such as a router or switch, in which some or all of the components **1010-1090** described herein are implemented using

application-specific integrated circuits (ASICs) or field programmable gate arrays (FPGAs). As another example, an implementing apparatus may include one or more memories storing instructions for implementing various components described herein, one or more hardware processors configured to execute the instructions stored in the one or more memories, and various data repositories in the one or more memories for storing data structures utilized and manipulated by various components **1010-1090**.

(219) Device **1000** is generally configured to receive and forward data units **1005** to other devices in a network, such as network **900**, by means of a series of operations performed at various components within the device **1000**. Note that certain nodes **910** in a system such as network **900** may each be or include a separate network device **1000**. In an embodiment, a node **910** may include more than one device **1000**. In an embodiment, device **1000** may itself be one of a number of components within a node **910**. For instance, network device **1000** may be an integrated circuit, or “chip,” dedicated to performing switching and/or routing functions within a network switch or router. The network switch or router may further comprise one or more central processor units, storage units, memories, physical interfaces, LED displays, or other components external to the chip, some or all of which may communicate with the chip.

(220) A non-limiting example flow of a data unit **1005** through various subcomponents of the switching logic of device **1000** is as follows. After being received via a port **1010**, a data unit **1005** may be buffered by an arbiter (not depicted) until the data unit **1005** can be processed by an ingress packet processor **1050**, and then delivered to an interconnect (not depicted). From the interconnect, the data unit **1005** may be forwarded to a traffic manager **1040**. The traffic manager **1040** may store the data unit **1005** in a buffer **1044** and assign the data unit **1005** to a queue **1045**. The traffic manager **1040** manages the flow of the data unit **1005** through the queue **1045** until the data unit **1005** is released to an egress packet processor **1050**. Depending on the processing, the traffic manager **1040** may then assign the data unit **1005** to another queue so that it may be processed by yet another egress processor **1050**, or the egress packet processor **1050** may send the data unit **1005** to an egress arbiter (not depicted) from which the data unit **1005** is finally forwarded out another port **1090**. Of course, depending on the embodiment, the switching logic may omit some of these subcomponents and/or include other subcomponents in varying arrangements.

(221) Example components of a device **1000** are now described in further detail.

(222) 5.3. Ports

(223) Network device **1000** includes ports **1010/1090**, by which the packet-handling logic is coupled to various network interfaces that receive and transmit data units **1005**. Ports **1010**, including ports **1010a-n**, are inbound (“ingress”) ports by which data units **1005** are received over a network, such as network **900**. Ports **1090**, including ports **1090a-n**, are outbound (“egress”) ports by which at least some of the data units **1005** are sent out to other destinations within the network, after having been processed by the network device **1000**.

(224) Data units **1005** may be of any suitable PDU type, such as packets, cells, frames, etc. In an embodiment, data units **1005** are packets. However, the individual atomic data units upon which the depicted components operate may actually be subunits of the data units **1005**. For example, at the port level, data units **1005** may be received, acted upon, and even transmitted as cells or other TDUs. These TDUs may be logically linked together as the data units **1005** (e.g. PDUs) to which they respectively belong for purposes of determining how to handle the cells or frames. However, the subunits may not actually be assembled into data units **1005** within device **1000**, particularly if the subunits are being forwarded to another destination through device **1000**.

(225) Ports **1010/1090** are depicted as separate ports for illustrative purposes, but may actually correspond to the same physical interface hardware on the network device **1010**. That is, a network device **1000** may both receive data units **1005** and send data units **1005** over a single physical interface, and the single physical interface may thus be logically coupled to both an ingress port **1010** and egress port **1090**. Nonetheless, for various functional purposes, certain logic of the

network device **1000** may view a single physical interface as a separate ingress port **1010** and egress port **1090**. Moreover, for various functional purposes, certain logic of the network device **1000** may subdivide a single physical interface into multiple ingress ports **1010** or egress ports **1090**, or aggregate multiple physical interfaces into a single ingress port **1010** or egress port **1090**. Hence, in various embodiments, ports **1010** and **1090** should be understood as distinct logical constructs that can be mapped to physical interfaces rather than simply as distinct physical constructs.

(226) In some embodiments, the physical interfaces to which the ports **1010/1090** of a device **1000** are coupled may include various interface components from a MAC layer and/or Physical layer, including one or more transceivers, Serializer/Deserializer (“SerDes”) blocks, interface controllers, and/or other suitable components, by which device **1000** receives, decodes, descrambles, deserializes, encodes, scrambles, serializes, sends, and/or otherwise processes data units on their way in or out of ports **1010/1090**.

(227) 5.4. Packet Processors

(228) A device **1000** comprises one or more packet processing components **1050**. These packet processors **1050** may be any suitable combination of fixed circuitry and/or software-based logic, such as specific logic components implemented by one or more Field Programmable Gate Arrays (FPGAs) or Application-Specific Integrated Circuits (ASICs), or a general-purpose processor executing software instructions.

(229) Different packet processors **1050** may be configured to perform different packet processing tasks. These tasks may include, for example, identifying paths along which to forward data units **1005**, forwarding data units **1005** to egress ports **1090**, implementing flow control and/or other policies, manipulating packets, performing statistical or debugging operations, and so forth. A device **1000** may comprise any number of packet processors **1050** configured to perform any number of processing tasks.

(230) In an embodiment, the packet processors **1050** within a device **1000** may be arranged such that the output of one packet processor **1050** may, eventually, be inputted into another packet processor **1050**, in such a manner as to pass data units **1005** from certain packet processor(s) **1050** to other packet processor(s) **1050** in a sequence of stages, until finally disposing of the data units **1005** (e.g. by sending the data units **1005** out an egress port **1090**, “dropping” the data units **1005**, etc.). The exact set and/or sequence of packet processors **1050** that process a given data unit **1005** may vary, in some embodiments, depending on attributes of the data unit **1005** and/or the state of the device **1000**. Any number of packet processors **1050** may be chained together in this manner.

(231) Based on decisions made while processing a data unit **1005**, a packet processor **1050** may, in some embodiments, and/or for certain processing tasks, manipulate a data unit **1005** directly. For instance, the packet processor **1050** may add, delete, or modify information in a data unit header or payload. In other embodiments, and/or for other processing tasks, a packet processor **1050** may generate control information that accompanies the data unit **1005**, or is merged with the data unit **1005**, as the data unit **1005** continues through the device **1000**. This control information may then be utilized by other components of the device **1000** to implement decisions made by the packet processor **1050**.

(232) In an embodiment, a packet processor **1050** need not necessarily process an entire data unit **1005**, but may rather only receive and process a subunit of a data unit **1005**, such as a TDU comprising header information for the data unit. For instance, if the data unit **1005** is a packet comprising multiple cells, the first cell, or a first subset of cells, might be forwarded to a packet processor **1050**, while the remaining cells of the packet (and potentially the first cell(s) as well) are forwarded in parallel to a merger component where they await results of the processing.

INGRESS AND EGRESS PROCESSORS

(233) In an embodiment, a packet processor may be generally classified as an ingress packet processor **1050** or an egress packet processor **1050**. Generally, an ingress processor **1050** resolves

destinations for a traffic manager **1040** to determine which ports **1090** and/or queues **1045** a data unit **1005** should depart from. There may be any number of ingress processors **1050**, including just a single ingress processor **1050**.

(234) In an embodiment, an ingress processor **1050** performs certain intake tasks on data units **1005** as they arrive. These intake tasks may include, for instance, and without limitation, parsing data units **1005**, performing routing related lookup operations, categorically blocking data units **1005** with certain attributes and/or when the device **1000** is in a certain state, duplicating certain types of data units **1005**, making initial categorizations of data units **1005**, and so forth. Once the appropriate intake task(s) have been performed, the data units **1005** are forwarded to an appropriate traffic manager **1040**, to which the ingress processor **1050** may be coupled directly or via various other components, such as an interconnect component.

(235) The egress packet processor(s) **1050** of a device **1000**, by contrast, may be configured to perform non-intake tasks necessary to implement the switching logic of the device **1000**. These tasks may include, for example, tasks such as identifying paths along which to forward the data units **1005**, implementing flow control and/or other policies, manipulating data units, performing statistical or debugging operations, and so forth. In an embodiment, there may be different egress packet processors(s) **1050** assigned to different flows or other categories of traffic, such that not all data units **1005** will be processed by the same egress packet processor **1050**.

(236) In an embodiment, each egress processor **1050** is coupled to a different group of egress ports **1090** to which they may send data units **1005** processed by the egress processor **1050**. In an embodiment, access to a group of ports **1090** may be regulated via an egress arbiter coupled to the egress packet processor **1050**. In some embodiments, an egress processor **1050** may also or instead be coupled to other potential destinations, such as an internal central processing unit, a storage subsystem, or a traffic manager **1040**.

(237) 5.5. Buffers

(238) Since not all data units **1005** received by the device **1000** can be processed by the packet processor(s) **1050** at the same time, various components of device **1000** may temporarily store data units **1005** in one or more buffers **1044** while the data units **1005** are waiting to be processed. For example, a certain packet processor **1050** may only be capable of processing a certain number of data units **1005**, or portions of data units **1005**, in a given clock cycle, meaning that other data units **1005**, or portions of data units **1005**, destined for the packet processor **1050** must either be ignored (e.g., dropped) or stored. At any given time, a large number of data units **1005** may be stored in the buffers **1044** of the device **1000**, depending on network traffic conditions.

(239) A device **1000** may include a variety of buffers **1044**, each utilized for varying purposes and/or components. Generally, a data unit **1005** awaiting processing by a component is held in a buffer **1044** associated with that component until the data unit **1005** is “released” to the component for processing. For example, a traffic manager **1040** will typically have a relatively large buffer **1044**, referred to as an egress buffer, in which it buffers data units **1005** prior to releasing those data units **1050** to an egress packet processor **1050**.

(240) A buffer **1044** may be implemented using a single physical memory unit (e.g. SRAM, DRAM, etc.), a designated portion of a memory unit, or a collection of memory units, depending on an embodiment. The buffer **1044** is divided into addressable units, or entries, that store SDUs, one or more of which form a TDU, as explained elsewhere. Each TDU stored in a buffer **1044** belongs to a PDU. However, the data for the SDUs and/or TDUs that belong to a PDU may not necessarily be stored adjacent to each other. If one wishes to reconstruct a PDU based on the buffered SDUs, one might be unable to do so using the buffer memory alone. Therefore, in an embodiment, a buffer **1044** may further store or be associated with linking data that indicates which SDUs belong to a given TDU and/or which TDUs belong to a given PDU, also referred to as intra-packet link data.

(241) For each PDU, buffer **1044** may further store or be associated with various PDU metadata.

The PDU metadata may include any suitable information about a PDU, such as a PDU identifier, location(s) of linking data for the PDU (e.g. the address(es) of intra-packet entr(ies) at which the linked list(s) for the PDU start), a count of TDUs in the PDU, source information, destination information, control information, timestamps, statistics, an assigned queue, flow control information, and so forth.

(242) 5.6. Queues

(243) In an embodiment, to manage the order in which data units **1005** are processed from buffers **1044**, various components of a device **1000** may implement queueing logic. For example, the flow of data units **1005** through the egress buffers **1044** of traffic manager **1040** may be managed using egress queues while the flow of data units **1005** through the buffers of an ingress arbiter might be managed using ingress queues.

(244) A queue **1045** is a set of nodes arranged in some order by metadata describing the queue **1045**. The queue **1045** includes a head node, or head, which is typically the next node to be processed, and a tail node, or tail, which is typically the node most recently added to the queue. A node will typically progress from the tail to the head over time as other nodes are processed and removed from the queue.

(245) In the case of queue **1045**, the nodes are data unit **1005**, or the buffer location(s) at which the data unit **1005** begins. A data unit **1005** that has been added to a queue **1045** is said to be “linked” to that queue **1045**. A data unit **1005** may belong to one or more queues **1045**.

(246) In many embodiments, the sequence in which the queue **1045** arranges its constituent data units **1005** generally corresponds to the order in which the data units **1005** or data unit portions in the queue **1045** will be released and processed. Such queues **1045** are known as first-in-first-out (“FIFO”) queues, though in other embodiments other types of queues may be utilized. In some embodiments, the number of data units **1005** or data unit portions assigned to a given queue **1045** at a given time may be limited, either globally or on a per-queue basis, and this limit may change over time.

(247) In an embodiment, queues **1045** are implemented using a linking memory referred to an “inter-packet” link memory, which is separate from the associated buffer memory **1044**. Each entry in the link memory is said to be a node in the queue. Each link entry points comprises a data pointer, which, when the link entry is occupied, points to a memory location in the buffer memory **1044** at which a corresponding data unit (or at least the start of the data unit) is found (e.g. a buffer entry, a first entry for the data unit in an intra-packet link memory, etc.). Each entry in the link memory further may further comprises a link pointer to another link entry, which corresponds to the next node in the queue. Of course, in other embodiments, other types of linking memories and/or other structures may instead be utilized instead to represent the queue.

(248) 5.7. Traffic Management

(249) According to an embodiment, a device **1000** further includes one or more traffic managers **1040** configured to control the flow of data units **1005** to one or more packet processor(s) **1050**. A traffic manager **1040** may receive data units **1005** directly from a port **1010**, from an ingress processor **1050**, and/or other suitable components of device **1000**. In an embodiment, the traffic manager **1040** is configured to receive up to one TDU from each possible source (e.g. each port **1010**) each clock cycle of the traffic manager **1040**.

(250) Traffic manager **1040** may include or be coupled to buffers **1044** for buffering data units **1005** prior to sending those data units **1005** to their respective processor(s) **1050**. A buffer manager within the traffic manager **1040** may temporarily store data units **1005** in buffers **1044** as they await processing by processor(s) **1050**. A data unit **1005** or data unit portion in a buffer **1044** may eventually be “released” to one or more processor(s) **1050** for processing, by reading the data unit **1005** from the buffer **1044** and sending the data unit **1005** to the processor(s) **1050**. In an embodiment, traffic manager **1040** may release up to a certain number of data units **1005** from buffers **1044** to processors **1050** each clock cycle.

(251) Beyond managing the use of buffers **1044** to store data units **1005** (or copies thereof), a traffic manager **1040** may include queue management logic configured to assign data units **1005** to queues **1045** and manage the flow of data units **1005** through queues **1045**. The traffic manager **1040** may, for instance, “enqueue” a PDU that has been fully buffered by identifying a specific queue **1045** to assign the PDU to, and then linking a PDU identifier or other PDU metadata to the assigned queue. The traffic manager **1040** may further determine when to release—also referred to as dequeuing—data units **1005** from queues **1045** by sending instructions to the buffer manager **1044** to read/release the data units **1005** and then providing the data read from the buffer **1044** to specific packet processor(s) **1050**.

(252) In an embodiment, different queues **1045** may exist for different sources or destinations. For example, each port **1010** and/or port **1090** may have its own set of queues **1045**. The queue **1045** to which an incoming data unit **1005** is assigned and linked may, for instance, be selected based on forwarding information indicating which port **1090** the data unit **1005** should depart from. In an embodiment, a different egress processor **1050** may be associated with each different set of one or more queues **1045**. In an embodiment, the current processing context of the data unit **1005** may be used to select which queue **1045** a data unit **1005** should be assigned to.

(253) In an embodiment, there may also or instead be different queues **1045** for different flows or sets of flows. That is, each identifiable traffic flow or group of traffic flows is assigned its own set of queues **1045** to which its data units **1005** are respectively assigned. In an embodiment, different queues **1045** may correspond to different classes of traffic or quality-of-service (QOS) levels. Different queues **1045** may also or instead exist for any other suitable distinguishing properties of the data units **1005**, such as source address, destination address, packet type, and so forth.

(254) Though only one packet processor **1050** and traffic manager **1040** are depicted, a device **1000** may comprise any number of packet processors **1050** and traffic managers **1040**. For instance, different sets of ports **1010** and/or ports **1090** may have their own traffic manager **1040** and packet processors **1050**. As another example, in an embodiment, the traffic manager **1040** may be duplicated for some or all of the stages of processing a data unit. For example, system **1000** may include a traffic manager **1040** and egress packet processor **1050** for an egress stage performed upon the data unit **1005** exiting the system **1000**, and/or a traffic manager **1040** and packet processor **1050** for any number of intermediate stages. The data unit **1005** may thus pass through any number of traffic managers **1040** and/or packet processors **1050** prior to exiting the system **1000**. In other embodiments, only a single traffic manager **1040** is needed. If intermediate processing is needed, flow of a data unit **1005** may loop back to the traffic manager **1040** for buffering and/or queuing after each stage of intermediate processing.

(255) In an embodiment, a traffic manager **1040** is coupled to the output of an ingress packet processor(s) **1050**, such that data units **1005** (or portions thereof) are assigned to buffers **1044** only upon being initially processed by an ingress packet processor **1050**. Once in an egress buffer **1044**, a data unit **1005** (or portion thereof) may be “released” to one or more egress packet processor(s) **1050** for processing.

(256) In the course of processing a data unit **1005**, a device **1000** may replicate a data unit **1005** one or more times for purposes such as, without limitation, multicasting, mirroring, debugging, and so forth. For example, a single data unit **1005** may be replicated to multiple egress queues **1045**. For instance, a data unit **1005** may be linked to separate queues **1045** for each of ports **1**, **3**, and **10**. As another example, a data unit **1005** may be replicated a number of times after it reaches the head of a queue **1045** (e.g. for different egress processors **1050**). Hence, though certain techniques described herein may refer to the original data unit **1005** that was received by the device **1000**, it will be understood that those techniques will equally apply to copies of the data unit **1005** that have been generated for various purposes. A copy of a data unit **1005** may be partial or complete. Moreover, there may be an actual physical copy of the data unit **1005** in buffers **1044**, or a single copy of the data unit **1005** may be linked from a single buffer location to multiple queues **1045** at the same

time.

(257) 5.8. Forwarding Logic

(258) The logic by which a device **1000** determines how to handle a data unit **1005**—such as where and whether to send a data unit **1005**, whether to perform additional processing on a data unit **1005**, etc. —is referred to as the forwarding logic of the device **1000**. This forwarding logic is collectively implemented by a variety of the components of the device **1000**, such as described above. For example, an ingress packet processor **1050** may be responsible for resolving the destination of a data unit **1005** and determining the set of actions/edits to perform on the data unit **1005**, and an egress packet processor **1050** may perform the edits. Also, there may be embodiments where the ingress packet processor **1050** performs edits as well.

(259) The forwarding logic may be hard-coded and/or configurable, depending on the embodiment. For example, the forwarding logic of a device **1000**, or portions thereof, may, in some instances, be at least partially hard-coded into one or more ingress processors **1050** and/or egress processors **1050**. As another example, the forwarding logic, or elements thereof, may also be configurable, in that the logic changes over time in response to analyses of state information collected from, or instructions received from, the various components of the device **1000** and/or other nodes in the network in which the device **1000** is located.

(260) In an embodiment, a device **1000** will typically store in its memories one or more forwarding tables (or equivalent structures) that map certain data unit attributes or characteristics to actions to be taken with respect to data units **1005** having those attributes or characteristics, such as sending a data unit **1005** to a selected path, or processing the data unit **1005** using a specified internal component. For instance, such attributes or characteristics may include a Quality-of-Service level specified by the data unit **1005** or associated with another characteristic of the data unit **1005**, a flow control group, an ingress port **1010** through which the data unit **1005** was received, a tag or label in a packet's header, a source address, a destination address, a packet type, or any other suitable distinguishing property. The forwarding logic may, for example, read such a table, determine one or more ports **1090** to send a data unit **1005** to based on the table, and add or associate the data unit **1005** with information that indicates to downstream logic to send the data unit **1005** along a path that includes a specific traffic manager **1040** and/or egress processor **1050** that is coupled to the one or more ports **1090**.

(261) According to an embodiment, the forwarding tables describe groups of one or more addresses, such as subnets of IPv4 or IPv6 addresses. Each address is an address of a network device on a network, though a network device may have more than one address. Each group is associated with a potentially different set of one or more actions to execute with respect to data units that resolve to (e.g. are directed to) an address within the group. Any suitable set of one or more actions may be associated with a group of addresses, including without limitation, forwarding a message to a specified “next hop,” duplicating the message, changing the destination of the message, dropping the message, performing debugging or statistical operations, applying a quality of service policy or flow control policy, and so forth.

(262) For illustrative purposes, these tables are described as “forwarding tables,” though it will be recognized that the extent of the action(s) described by the tables may be much greater than simply where to forward the message. For example, in an embodiment, a table may be a basic forwarding table that simply specifies a next hop for each group. In other embodiments, a table may describe one or more complex policies for each group. Moreover, there may be different types of tables for different purposes. For instance, one table may be a basic forwarding table that is compared to the destination address of each packet, while another table may specify policies to apply to packets upon ingress based on their destination (or source) group, and so forth.

(263) In an embodiment, forwarding logic may read port state data for ports **1010/1090**. Port state data may include, for instance, flow control state information describing various traffic flows and associated traffic flow control rules or policies, link status information indicating links that are up

or down, port utilization information indicating how ports are being utilized (e.g. utilization percentages, utilization states, etc.). Forwarding logic may be configured to implement the associated rules or policies associated with the flow(s) to which a given packet belongs.

(264) As data units **1005** are routed through different nodes in a network, the nodes may, on occasion, discard, fail to send, or fail to receive certain data units **1005**, thus resulting in the data units **1005** failing to reach their intended destination. The act of discarding of a data unit **1005**, or failing to deliver a data unit **1005**, is typically referred to as “dropping” the data unit. Instances of dropping a data unit **1005**, referred to herein as “drops” or “packet loss,” may occur for a variety of reasons, such as resource limitations, errors, or deliberate policies. Different components of a device **1000** may make the decision to drop a data unit **1005** for various reasons. For instance, a traffic manager **1040** may determine to drop a data unit **1005** because, among other reasons, buffers **1044** are overutilized, a queue **1045** is over a certain size, and/or a data unit **1005** has a certain characteristic.

(265) 5.9. Multi-Pipeline Architecture

(266) In an embodiment, a network device may include multiple pipelines of data unit processing components such as those described above. FIG. **11** illustrates an example of one such network device with multiple packet processing pipelines, according to an embodiment. Network device **1100** includes a plurality of ingress ports **1110** and egress ports **1190**, similar to the ingress ports **1010** and egress ports **1090** of device **1000**. The ingress ports **1110** are divided into port groups **1110a-n**, and each group of ports **1110** feeds data units to a different pipeline **1102** of processing components. There may be any number of groups of ports **1110**, and hence any number of corresponding pipelines **1102**.

(267) Each pipeline includes an ingress arbiter **1120**. Each ingress arbiter **1120** is coupled to a corresponding group of ingress ports **1110**, and receives data units from those ports **1110**. In some respects, each ingress arbiter **1120** may be viewed as an ingress version of traffic manager **1040**. An ingress arbiter **1120** is responsible for determining when data units are sent to downstream components, and in particular to an ingress packet processor **1130** that is coupled to the ingress arbiter **1120** within a corresponding pipeline **1102**. An ingress arbiter **1120** may or may not include its own buffer memory in which it buffers data unit that await processing, depending on the embodiment.

(268) In an embodiment, the data units sent by the ingress arbiter **1120** are actually subunits, such as cells, frames, segments, or other TDUs, of larger parent data units, such as packets or other PDUs. The ingress arbiter **1120** may, in an embodiment, divide a larger data unit (or ensure that the larger data unit is divided) into these subunits prior to transmitting the data units to the corresponding ingress packet processor **1130**. In an embodiment, a packet or other PDU may arrive at the ingress arbiter **1120** as a set of TDUs. For convenience, examples are given herein where the TDU is a cell, and the PDU is a packet, but it will be appreciated that the cell may in fact be any type of subunit, and the packet may in fact be any larger data unit that comprises those subunits.

(269) Each pipeline **1102** further includes an ingress packet processor **1130** to which its ingress arbiter **1120** eventually sends these data units. Each ingress packet processor **1130**, meanwhile, functions in similar manner as an ingress packet processor **1050** described above. In particular, an ingress packet processor **1130** performs forwarding tasks such as resolving the data unit destination, adding or removing headers, and so forth. For instance, the ingress packet processor **1130** may be responsible for generating control information that instructs downstream components of the pipelines **1102** on how to handle the data unit, and this control information may either be inserted into the data unit, or be conveyed along with the data unit as sideband information.

(270) Each pipeline **1102** further includes an egress traffic manager **1140**, which functions in similar manner to the traffic manager **1040**. A common interconnect **1138** is coupled to each ingress packet processor **1130** on one end and each egress traffic manager **1140** on the other. The interconnect **1138** conveys data units to traffic manager(s) **1140** indicated by the ingress packet

processors **1130** (e.g. based on the control information), thereby allowing data units to “switch” pipelines **1102** should their destination(s) include a port **1190** that is not on the same pipeline **1102** as the ingress port **1110** through which the data unit was received. Ingress packet processors **1130** may be coupled to the interconnect **1138** directly, or indirectly via other components such as a merger unit (e.g. that merges a control portion of the data unit processed by the packet processor **1130** with a data portion of the data unit that bypasses the packet processor **1130**).

(271) A pipeline's egress traffic manager **1140** then regulates the flow of data units to the pipeline's egress packet processor **1150**, in similar manner as described with respect to traffic manager **1040**. The egress packet processor **1150** processes data units in similar manner as described with respect to egress packet processors **1050**. The egress packet processors then forward the processed data units to the pipeline's egress port transmit unit **1160**, which is responsible for transmission of data units out a set of one or more egress ports **1190** belonging to the pipeline **1102**. The set of egress ports **1190** for a pipeline corresponds to the pipeline's ingress ports **1110**.

(272) In yet other embodiments, an egress traffic manager **1140** and/or other components may be shared between such pipelines.

(273) 5.10. Miscellaneous

(274) Devices **1000** and **1100** illustrate only several of many possible arrangements of components in a network device configured to provide the functionality described herein. Other arrangements may include fewer, additional, or different components, and the division of work between the components may vary depending on the arrangement. Moreover, in an embodiment, the techniques described herein may be utilized in a variety of computing contexts other than within a network **900**.

(275) For simplification, the traffic managers, packet processors, and other components are on occasion described herein as acting upon or with respect to a data unit, when in fact only a portion of the data unit may be sent to or otherwise available to the component. For instance, a packet processor may be said to process a data unit, when in fact only the data unit control portion is available to the packet processor. In such contexts, it will be understood that the information about the data unit that is actually available to the component is a representation of the data unit to that component. Any actions described herein as having been taken by the component on or with respect to the data unit may occur with respect to the representation of the data unit, and not necessarily the complete data unit.

(276) It will be appreciated that the actual physical representation of a data unit may change as a result of the processes described herein. For instance, a data unit may be converted from a physical representation at a particular location in one memory to a signal-based representation, and back to a physical representation at a different location in a potentially different memory, as it is moved from one component to another within a network device or even between network devices. Such movement may technically involve deleting, converting, and/or copying some or all of the data unit any number of times. For simplification, however, the data unit is logically said to remain the same data unit as it moves through the device, even if the physical representation of the data unit changes. Similarly, the contents and/or structure of a data unit may change as it is processed, such as by adding or deleting header information, adjusting cell boundaries, or even modifying payload data. A modified data unit is nonetheless still said to be the same data unit, even after altering its contents and/or structure.

6.0. Implementation Mechanism—Hardware Overview

(277) According to one embodiment, the techniques described herein are implemented by one or more special-purpose computing devices. The special-purpose computing devices may be desktop computer systems, portable computer systems, handheld devices, networking devices, or any other device that incorporates hard-wired and/or program logic to implement the techniques. The special-purpose computing devices may be hard-wired to perform the techniques, or may include digital electronic devices such as one or more application-specific integrated circuits (ASICs) or field

programmable gate arrays (FPGAs) that are persistently programmed to perform the techniques. Such special-purpose computing devices may also combine custom hard-wired logic, ASICs, FPGAs, or other circuitry with custom programming to accomplish the techniques.

(278) Though certain foregoing techniques are described with respect to a hardware implementation, which provides a number of advantages in certain embodiments, it will also be recognized that, in other embodiments, the foregoing techniques may still provide certain advantages when performed partially or wholly in software. Accordingly, in such an embodiment, a suitable implementing apparatus comprises a general-purpose hardware processor and is configured to perform any of the foregoing methods by executing program instructions in firmware, memory, other storage, or a combination thereof.

(279) FIG. 12 is a block diagram that illustrates an example computer system 1200 that may be utilized in implementing the above-described techniques, according to an embodiment. Computer system 1200 may be, for example, a desktop computing device, laptop computing device, tablet, smartphone, server appliance, computing mainframe, multimedia device, handheld device, networking apparatus, or any other suitable device. In an embodiment, FIG. 12 constitutes a different view of the devices and systems described in previous sections.

(280) Computer system 1200 may include one or more ASICs, FPGAs, or other specialized circuitry 1203 for implementing program logic as described herein. For example, circuitry 1203 may include fixed and/or configurable hardware logic blocks for implementing some or all of the described techniques, input/output (I/O) blocks, hardware registers or other embedded memory resources such as random-access memory (RAM) for storing various data, and so forth. The logic blocks may include, for example, arrangements of logic gates, flip-flops, multiplexers, and so forth, configured to generate an output signals based on logic operations performed on input signals.

(281) Additionally, and/or instead, computer system 1200 may include one or more hardware processors 1204 configured to execute software-based instructions. Computer system 1200 may also include one or more busses 1202 or other communication mechanism for communicating information. Busses 1202 may include various internal and/or external components, including, without limitation, internal processor or memory busses, a Serial ATA bus, a PCI Express bus, a Universal Serial Bus, a HyperTransport bus, an InfiniBand bus, and/or any other suitable wired or wireless communication channel.

(282) Computer system 1200 also includes one or more memories 1206, such as a RAM, hardware registers, or other dynamic or volatile storage device for storing data units to be processed by the one or more ASICs, FPGAs, or other specialized circuitry 1203. Memory 1206 may also or instead be used for storing information and instructions to be executed by processor 1204. Memory 1206 may be directly connected or embedded within circuitry 1203 or a processor 1204. Or, memory 1206 may be coupled to and accessed via bus 1202. Memory 1206 also may be used for storing temporary variables, data units describing rules or policies, or other intermediate information during execution of program logic or instructions.

(283) Computer system 1200 further includes one or more read only memories (ROM) 1208 or other static storage devices coupled to bus 1202 for storing static information and instructions for processor 1204. One or more storage devices 1210, such as a solid-state drive (SSD), magnetic disk, optical disk, or other suitable non-volatile storage device, may optionally be provided and coupled to bus 1202 for storing information and instructions.

(284) A computer system 1200 may also include, in an embodiment, one or more communication interfaces 1218 coupled to bus 1202. A communication interface 1218 provides a data communication coupling, typically two-way, to a network link 1220 that is connected to a local network 1222. For example, a communication interface 1218 may be an integrated services digital network (ISDN) card, cable modem, satellite modem, or a modem to provide a data communication connection to a corresponding type of telephone line. As another example, the one or more communication interfaces 1218 may include a local area network (LAN) card to provide a data

communication connection to a compatible LAN. As yet another example, the one or more communication interfaces **1218** may include a wireless network interface controller, such as an 802.11-based controller, Bluetooth controller, Long Term Evolution (LTE) modem, and/or other types of wireless interfaces. In any such implementation, communication interface **1218** sends and receives electrical, electromagnetic, or optical signals that carry digital data streams representing various types of information.

(285) Network link **1220** typically provides data communication through one or more networks to other data devices. For example, network link **1220** may provide a connection through local network **1222** to a host computer **1224** or to data equipment operated by a Service Provider **1226**. Service Provider **1226**, which may for example be an Internet Service Provider (ISP), in turn provides data communication services through a wide area network, such as the world-wide packet data communication network now commonly referred to as the “Internet” **1228**. Local network **1222** and Internet **1228** both use electrical, electromagnetic or optical signals that carry digital data streams. The signals through the various networks and the signals on network link **1220** and through communication interface **1218**, which carry the digital data to and from computer system **1200**, are example forms of transmission media.

(286) In an embodiment, computer system **1200** may send and receive data units through the network(s), network link **1220**, and communication interface **1218**. In some embodiments, this data may be data units that the computer system **1200** has been asked to process and, if necessary, redirect to other computer systems via a suitable network link **1220**. In other embodiments, this data may be instructions for implementing various processes related to the described techniques. For instance, in the Internet example, a server **1230** might transmit a requested code for an application program through Internet **1228**, ISP **1226**, local network **1222** and communication interface **1218**. The received code may be executed by processor **1204** as it is received, and/or stored in storage device **1210**, or other non-volatile storage for later execution. As another example, information received via a network link **1220** may be interpreted and/or processed by a software component of the computer system **1200**, such as a web browser, application, or server, which in turn issues instructions based thereon to a processor **1204**, possibly via an operating system and/or other intermediate layers of software components.

(287) Computer system **1200** may optionally be coupled via bus **1202** to one or more displays **1212** for presenting information to a computer user. For instance, computer system **1200** may be connected via a High-Definition Multimedia Interface (HDMI) cable or other suitable cabling to a Liquid Crystal Display (LCD) monitor, and/or via a wireless connection such as peer-to-peer Wi-Fi Direct connection to a Light-Emitting Diode (LED) television. Other examples of suitable types of displays **1212** may include, without limitation, plasma display devices, projectors, cathode ray tube (CRT) monitors, electronic paper, virtual reality headsets, braille terminal, and/or any other suitable device for outputting information to a computer user. In an embodiment, any suitable type of output device, such as, for instance, an audio speaker or printer, may be utilized instead of a display **1212**.

(288) One or more input devices **1214** are optionally coupled to bus **1202** for communicating information and command selections to processor **1204**. One example of an input device **1214** is a keyboard, including alphanumeric and other keys. Another type of user input device **1214** is cursor control **1216**, such as a mouse, a trackball, or cursor direction keys for communicating direction information and command selections to processor **1204** and for controlling cursor movement on display **1212**. This input device typically has two degrees of freedom in two axes, a first axis (e.g., x) and a second axis (e.g., y), that allows the device to specify positions in a plane. Yet other examples of suitable input devices **1214** include a touch-screen panel affixed to a display **1212**, cameras, microphones, accelerometers, motion detectors, and/or other sensors. In an embodiment, a network-based input device **1214** may be utilized. In such an embodiment, user input and/or other information or commands may be relayed via routers and/or switches on a Local Area Network (LAN) or other suitable shared network, or via a peer-to-peer network, from the input device **1214**

to a network link **1220** on the computer system **1200**.

(289) As discussed, computer system **1200** may implement techniques described herein using customized hard-wired logic, one or more ASICs or FPGAs **1203**, firmware and/or program logic, which in combination with the computer system causes or programs computer system **1200** to be a special-purpose machine. According to one embodiment, however, the techniques herein are performed by computer system **1200** in response to processor **1204** executing one or more sequences of one or more instructions contained in main memory **1206**. Such instructions may be read into main memory **1206** from another storage medium, such as storage device **1210**. Execution of the sequences of instructions contained in main memory **1206** causes processor **1204** to perform the process steps described herein.

(290) The term “storage media” as used herein refers to any non-transitory media that store data and/or instructions that cause a machine to operate in a specific fashion. Such storage media may comprise non-volatile media and/or volatile media. Non-volatile media includes, for example, optical or magnetic disks, such as storage device **1210**. Volatile media includes dynamic memory, such as main memory **1206**. Common forms of storage media include, for example, a floppy disk, a flexible disk, hard disk, solid state drive, magnetic tape, or any other magnetic data storage medium, a CD-ROM, any other optical data storage medium, any physical medium with patterns of holes, a RAM, a PROM, and EPROM, a FLASH-EPROM, NVRAM, any other memory chip or cartridge.

(291) Storage media is distinct from but may be used in conjunction with transmission media. Transmission media participates in transferring information between storage media. For example, transmission media includes coaxial cables, copper wire and fiber optics, including the wires that comprise bus **1202**. Transmission media may also take the form of acoustic or light waves, such as those generated during radio-wave and infra-red data communications.

(292) Various forms of media may be involved in carrying one or more sequences of one or more instructions to processor **1204** for execution. For example, the instructions may initially be carried on a magnetic disk or solid-state drive of a remote computer. The remote computer can load the instructions into its dynamic memory and use a modem to send the instructions over a network, such as a cable network or cellular network, as modulated signals. A modem local to computer system **1200** may receive the data on the network and demodulate the signal to decode the transmitted instructions. Appropriate circuitry may then place the data on bus **1202**. Bus **1202** carries the data to main memory **1206**, from which processor **1204** retrieves and executes the instructions. The instructions received by main memory **1206** may optionally be stored on storage device **1210** either before or after execution by processor **1204**.

7.0. Extensions and Alternatives

(293) As used herein, the terms “first,” “second,” “certain,” and “particular” are used as naming conventions to distinguish queries, plans, representations, steps, objects, devices, or other items from each other, so that these items may be referenced after they have been introduced. Unless otherwise specified herein, the use of these terms does not imply an ordering, timing, or any other characteristic of the referenced items.

(294) In the drawings, the various components are depicted as being communicatively coupled to various other components by arrows. These arrows illustrate only certain examples of information flows between the components. Neither the direction of the arrows nor the lack of arrow lines between certain components should be interpreted as indicating the existence or absence of communication between the certain components themselves. Indeed, each component may feature a suitable communication interface by which the component may become communicatively coupled to other components as needed to accomplish any of the functions described herein.

(295) In the foregoing specification, embodiments of the inventive subject matter have been described with reference to numerous specific details that may vary from implementation to implementation. Thus, the sole and exclusive indicator of what is the inventive subject matter, and

is intended to be the inventive subject matter, is the set of claims that issue from this application, in the specific form in which such claims issue, including any subsequent correction. In this regard, although specific claim dependencies are set out in the claims of this application, it is to be noted that the features of the dependent claims of this application may be combined as appropriate with the features of other dependent claims and with the features of the independent claims of this application, and not merely according to the specific dependencies recited in the set of claims. Moreover, although separate embodiments are discussed herein, any combination of embodiments and/or partial embodiments discussed herein may be combined to form further embodiments. (296) Any definitions expressly set forth herein for terms contained in such claims shall govern the meaning of such terms as used in the claims. Hence, no limitation, element, property, feature, advantage or attribute that is not expressly recited in a claim should limit the scope of such claim in any way. The specification and drawings are, accordingly, to be regarded in an illustrative rather than a restrictive sense.

Claims

1. A network apparatus comprising: a plurality of ports configured to send and/or receive data units over one or more networks; enqueue logic configured to enqueue the data units in queues, in which the data units await transmission to one or more downstream logic components, wherein each queue of the queues is assigned to one of the ports, the queue holding data units that are to be sent to or were received from the port assigned to the queue; dequeue logic configured to dequeue particular data units from the queues in response to instructions from a scheduler, wherein when instructed by the scheduler to dequeue a data unit associated with a particular port, selects a particular queue, of the queues, assigned to the particular port from which to dequeue a particular data unit at the head of the particular queue; a data unit transmitter configured to send the particular data units that have been dequeued by the dequeue logic to the one or more downstream logic components; the scheduler, configured to determine, at a given time, a given port of the ports to service, and send an instruction to the dequeue logic to dequeue data associated with the given port at the given time, the scheduler configured to service at least a first port of the ports no more frequently than at a service rate, the service rate indicates how frequently a port may be serviced in a clock cycle; a rate adjuster configured to, based on sizes of certain data units, of the data units, that are assigned to the first port, temporarily adjust the service rate from a first rate to a second rate.
2. The apparatus of claim 1, wherein the rate adjuster is configured to adjust the service rate from the first rate to the second rate responsive to a determination that at least two data units associated with the first port are both ready to be dequeued and are smaller than a particular size.
3. The apparatus of claim 2, wherein determining that at least two data units associated with the first port are smaller than the particular size comprises at least one of: determining that the at least two data units are, in combination, smaller than a maximum data unit size.
4. The apparatus of claim 1, wherein the rate adjuster is configured to adjust the service rate from the first rate to the second rate responsive to a determination that an average size of a plurality of the data units over a recent period of time is smaller than a particular size.
5. The apparatus of claim 1, wherein the first rate is a rate at which the first port must be serviced to ensure that a defined bandwidth of the first port is realized when the first port is transmitting maximum-size data units, wherein the second rate is greater than the first rate.
6. The apparatus of claim 1, wherein the rate adjuster is configured to adjust the service rate by at least causing the dequeue logic to dequeue at least two data units for the first port concurrently at a particular time.
7. The apparatus of claim 1, wherein the rate adjuster is configured to adjust the service rate by at least causing the enqueue logic to link two or more data units together when enqueueing the two or

more data units in a queue associated with the first port; wherein the dequeue logic is configured to dequeue the two or more data units from the queue concurrently at a particular time in a manner that is transparent to the dequeue logic as a consequence of the linking.

8. The apparatus of claim 1, wherein the rate adjuster is configured to adjust the service rate by at least causing the dequeue logic to dequeue multiple data units for the first port at a same time, wherein the data unit transmitter is configured to send one data unit for the first port to the one or more downstream logic components at a time, wherein the data unit transmitter includes a port FIFO for the first port, the port FIFO configured to store at least one of the multiple data units for transmission to the one or more downstream logic components at a later time.

9. The apparatus of claim 8, wherein, when the port FIFO stores at least one data unit, the rate adjuster is configured to activate a pseudo-port for the first port in a port scheduling sequence for the scheduler, the first port and the pseudo-port separately selectable from the sequence at a frequency corresponding to the first rate, the first port thereby effectively serviced at the second rate.

10. The apparatus of claim 8, wherein, when the port FIFO stores at least one data unit, and the scheduler determines whether to service the first port or a corresponding pseudo-port in a given clock cycle, the scheduler is configured to instruct the transmission logic to transmit a next data unit in the port FIFO to the one or more downstream logic components instead of a newly dequeued data unit.

11. The apparatus of claim 1, wherein, when a particular data unit is an end-of-packet data unit and the rate adjuster detects another end-of-packet data unit at the head of another queue assigned to the particular port, or immediately following the particular data unit in the particular queue, the dequeue logic is configured to dequeue both the particular data unit and the another end-of-packet data unit at the same time.

12. The apparatus of claim 1, wherein the one or more downstream logic components being one or more traffic managers, one or more packet processors, or the plurality of ports.

13. A method comprising: sending and/or receiving data units over one or more networks via a plurality of ports of a network apparatus; enqueueing the data units in queues, in which the data units await transmission to one or more downstream logic components of the network apparatus, wherein each queue of the data unit queues is assigned to one of the ports, the queue holding data units that are to be sent to or were received from the port assigned to the queue; dequeuing particular data units from the queues, the dequeuing including determining, at a given time, a given port of the ports to service, based at least on a constraint that a first port of the ports is serviced no more frequently than at a service rate, the service rate indicates how frequently a port may be serviced in a clock cycle, and dequeuing data associated with the given port at the given time, wherein when dequeuing data associated with a particular port, selecting a particular queue, of the queues, assigned to the particular port from which to dequeue a particular data unit at the head of the particular queue; sending the particular data units that have been dequeued to the one or more downstream logic components; based on sizes of certain data units, of the data units, that are assigned to the first port, temporarily adjusting the service rate from a first rate to a second rate.

14. The method of claim 13, wherein adjusting the service rate from the first rate to the second rate is responsive to a determination that at least two data units associated with the first port are both ready to be dequeued and are smaller than a particular size.

15. The method of claim 13, wherein the first rate is a rate at which the first port must be serviced to ensure that a defined bandwidth of the first port is realized when the first port is transmitting maximum-size data units, wherein the second rate is greater than the first rate.

16. The method of claim 13, wherein adjusting the service rate comprises dequeuing the at least two data units concurrently at a particular time.

17. The method of claim 13, wherein adjusting the service rate comprises dequeuing multiple data units for the first port at a same time, wherein only one data unit is sent for the first port to the one

or more downstream logic components at a time, wherein the method further comprises: storing at least one of the multiple data units in a port FIFO for the first port; activating a pseudo-port for the first port in a port scheduling sequence when the port FIFO stores at least one data unit; wherein determining the given port to service at the given time comprises selecting the given port using the port scheduling sequence, the first port and the pseudo-port separately selectable from the sequence at a frequency corresponding to the first rate, the first port thereby effectively serviced at the second rate; and wherein when the port FIFO stores at least one data unit, upon determining to service the first port or a corresponding pseudo-port in a given clock cycle, transmitting a next data unit from the port FIFO to the one or more downstream logic components, instead of a newly dequeued data unit.

18. The method of claim 13, further comprising: dequeuing both a particular data unit and another end-of-packet data unit at a same time when the particular data unit is an end-of-packet data unit and the another end-of-packet data unit is at the head of another queue assigned to the particular port, or immediately follows the particular data unit in the particular queue.

19. The method of claim 13, wherein the one or more downstream logic components is one or more traffic managers, one or more packet processors, or the plurality of ports.

20. The method of claim 13, wherein adjusting the service rate from the first rate to the second rate is responsive to a determination that an average size of a plurality of the data units over a recent period of time is smaller than a particular size.
