



US012413538B2

(12) **United States Patent**
Mitulal et al.

(10) **Patent No.:** **US 12,413,538 B2**

(45) **Date of Patent:** ***Sep. 9, 2025**

(54) **QUEUE PROTECTION USING A SHARED GLOBAL MEMORY RESERVE**

(71) Applicant: **Cisco Technology, Inc.**, San Jose, CA (US)

(72) Inventors: **Vinod Mitulal**, Santa Clara, CA (US); **Krishnan Subramani**, San Jose, CA (US); **Peter Newman**, Fremont, CA (US); **Georges Akis**, Los Altos, CA (US)

(73) Assignee: **Cisco Technology, Inc.**, San Jose, CA (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 13 days.

This patent is subject to a terminal disclaimer.

(21) Appl. No.: **18/153,923**

(22) Filed: **Jan. 12, 2023**

(65) **Prior Publication Data**

US 2023/0145162 A1 May 11, 2023

Related U.S. Application Data

(63) Continuation of application No. 15/250,860, filed on Aug. 29, 2016, now Pat. No. 11,563,695.

(51) **Int. Cl.**
H04L 49/103 (2022.01)
H04L 47/10 (2022.01)
(Continued)

(52) **U.S. Cl.**
CPC **H04L 49/103** (2013.01); **H04L 47/29** (2013.01); **H04L 47/30** (2013.01); **H04L 47/32** (2013.01);
(Continued)

(58) **Field of Classification Search**

CPC H04L 47/00–829; H04L 49/103; H04L 49/108; H04L 49/3036; H04L 49/9005; H04L 49/9084

See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

4,688,695 A 8/1987 Hirohata
5,263,003 A 11/1993 Cowles et al.
(Continued)

FOREIGN PATENT DOCUMENTS

CN 1777147 A 5/2006
EP 2228719 A1 9/2010
(Continued)

OTHER PUBLICATIONS

International Search Report and Written Opinion from the International Searching Authority for the corresponding International Application No. PCT/US2017/043463, mailed Oct. 16, 2017, 11 pages.
(Continued)

Primary Examiner — Michael Thier


Assistant Examiner — Prince A Mensah

(74) *Attorney, Agent, or Firm* — Polsinelli P.C.

(57) **ABSTRACT**

The subject technology relates to the management of a shared buffer memory in a network switch. Systems, methods, and machine readable media are provided for receiving a data packet at a first network queue from among a plurality of network queues, determining if a fill level of a queue in a shared buffer of the network switch exceeds a dynamic queue threshold, and in an event that the fill level of the shared buffer exceeds the dynamic queue threshold, determining if a fill level of the first network queue is less than a static queue minimum threshold.

20 Claims, 5 Drawing Sheets



(51)	Int. Cl.		6,907,419 B1	6/2005	Pesola et al.
	H04L 47/30	(2022.01)	6,912,668 B1	6/2005	Brown et al.
	H04L 47/32	(2022.01)	6,952,734 B1	10/2005	Gunlock et al.
	H04L 49/00	(2022.01)	6,976,090 B2	12/2005	Ben-Shaul et al.
	H04L 49/104	(2022.01)	6,978,300 B1	12/2005	Beukema et al.
	H04L 49/90	(2022.01)	6,983,303 B2	1/2006	Pellegrino et al.
(52)	U.S. Cl.		6,986,015 B2	1/2006	Testardi
	H04L 49/9005	(2022.01)	6,986,069 B2	1/2006	Oehler et al.
	CPC	H04L 49/108 (2013.01); H04L 49/3036	7,051,056 B2	5/2006	Rodriguez-Rivera et al.
		(2013.01); H04L 49/9005 (2013.01); H04L	7,069,465 B2	6/2006	Chu et al.
		49/9084 (2013.01)	7,073,017 B2	7/2006	Yamamoto
			7,108,339 B2	9/2006	Berger
(56)	References Cited		7,149,858 B1	12/2006	Kiselev
	U.S. PATENT DOCUMENTS		7,171,514 B2	1/2007	Coronado et al.
	5,339,445 A	8/1994 Gasztonyi	7,171,668 B2	1/2007	Molloy et al.
	5,430,859 A	7/1995 Norman et al.	7,174,354 B2	2/2007	Andreasson
	5,457,746 A	10/1995 Dolphin	7,200,144 B2	4/2007	Terrell et al.
	5,535,336 A	7/1996 Smith et al.	7,222,255 B1	5/2007	Claessens et al.
	5,588,012 A	12/1996 Oizumi	7,237,045 B2	6/2007	Beckmann et al.
	5,617,421 A	4/1997 Chin et al.	7,240,188 B2	7/2007	Takata et al.
	5,680,579 A	10/1997 Young et al.	7,246,260 B2	7/2007	Brown et al.
	5,690,194 A	11/1997 Parker et al.	7,266,718 B2	9/2007	Idei et al.
	5,740,171 A	4/1998 Mazzola et al.	7,269,168 B2	9/2007	Roy et al.
	5,742,604 A	4/1998 Edsall et al.	7,277,431 B2	10/2007	Walter et al.
	5,764,636 A	6/1998 Edsall	7,277,948 B2	10/2007	Igarashi et al.
	5,809,285 A	9/1998 Hilland	7,305,658 B1	12/2007	Hamilton et al.
	5,812,814 A	9/1998 Sukegawa	7,328,434 B2	2/2008	Swanson et al.
	5,812,950 A	9/1998 Tom	7,340,555 B2	3/2008	Ashmore et al.
	5,838,970 A	11/1998 Thomas	7,346,751 B2	3/2008	Prahlad et al.
	5,999,930 A	12/1999 Wolff	7,352,706 B2	4/2008	Klotz et al.
	6,035,105 A	3/2000 McCloghrie et al.	7,353,305 B2	4/2008	Pangal et al.
	6,043,777 A	3/2000 Bergman et al.	7,359,321 B1 *	4/2008	Sindhu H04L 47/11 710/52
	6,101,497 A	8/2000 Ofek	7,383,381 B1	6/2008	Faulkner et al.
	6,148,414 A	11/2000 Brown et al.	7,403,987 B1	7/2008	Marinelli et al.
	6,185,203 B1	2/2001 Berman	7,433,326 B2	10/2008	Desai et al.
	6,188,694 B1	2/2001 Fine et al.	7,433,948 B2	10/2008	Edsall et al.
	6,202,135 B1	3/2001 Kedem et al.	7,434,105 B1	10/2008	Rodriguez-Rivera et al.
	6,208,649 B1	3/2001 Kloth	7,441,154 B2	10/2008	Klotz et al.
	6,209,059 B1	3/2001 Ofer et al.	7,447,839 B2	11/2008	Uppala
	6,219,699 B1	4/2001 McCloghrie et al.	7,487,321 B2	2/2009	Muthiah et al.
	6,219,753 B1	4/2001 Richardson	7,500,053 B1	3/2009	Kavuri et al.
	6,223,250 B1	4/2001 Yokono	7,512,744 B2	3/2009	Banga et al.
	6,226,771 B1	5/2001 Hilla et al.	7,542,681 B2	6/2009	Cornell et al.
	6,260,120 B1	7/2001 Blumenau et al.	7,558,872 B1	7/2009	Senevirathne et al.
	6,266,705 B1	7/2001 Ullum et al.	7,587,570 B2	9/2009	Sarkar et al.
	6,269,381 B1	7/2001 St. Pierre et al.	7,631,023 B1	12/2009	Kaiser et al.
	6,269,431 B1	7/2001 Dunham	7,643,505 B1	1/2010	Colloff
	6,295,575 B1	9/2001 Blumenau et al.	7,654,625 B2	2/2010	Amann et al.
	6,400,730 B1	6/2002 Latif et al.	7,657,796 B1	2/2010	Kaiser et al.
	6,408,406 B1	6/2002 Parris	7,668,981 B1	2/2010	Nagineni et al.
	6,456,590 B1	9/2002 Ren et al.	7,669,071 B2	2/2010	Cochran et al.
	6,539,024 B1 *	3/2003 Janoska H04L 47/24 370/413	7,689,384 B1	3/2010	Becker
	6,542,909 B1	4/2003 Tamer et al.	7,694,092 B2	4/2010	Mizuno
	6,542,961 B1	4/2003 Matsunami et al.	7,697,554 B1	4/2010	Ofer et al.
	6,553,390 B1	4/2003 Gross et al.	7,706,303 B2	4/2010	Bose et al.
	6,564,252 B1	5/2003 Hickman et al.	7,707,481 B2	4/2010	Kirschner et al.
	6,647,474 B2	11/2003 Yanai et al.	7,716,648 B2	5/2010	Vaidyanathan et al.
	6,674,718 B1 *	1/2004 Heddes H04L 47/29 710/52	7,752,360 B2	7/2010	Galles
	6,675,258 B1	1/2004 Bramhall et al.	7,757,059 B1	7/2010	Ofer et al.
	6,683,883 B1	1/2004 Czeiger et al.	7,774,329 B1	8/2010	Peddy et al.
	6,694,413 B1	2/2004 Mimatsu et al.	7,774,839 B2	8/2010	Nazzari
	6,708,227 B1	3/2004 Cabrera et al.	7,793,138 B2	9/2010	Rastogi et al.
	6,715,007 B1	3/2004 Williams et al.	7,840,730 B2	11/2010	D'Amato et al.
	6,728,791 B1	4/2004 Young	7,843,906 B1	11/2010	Chidambaram et al.
	6,772,231 B2	8/2004 Reuter et al.	7,895,428 B2	2/2011	Boland, IV et al.
	6,820,099 B1	11/2004 Huber et al.	7,904,599 B1	3/2011	Bennett
	6,847,647 B1	1/2005 Wrenn	7,930,494 B1	4/2011	Goheer et al.
	6,848,759 B2	2/2005 Doornbos et al.	7,975,175 B2	7/2011	Votta et al.
	6,850,955 B2	2/2005 Sonoda et al.	7,979,670 B2	7/2011	Saliba et al.
	6,876,656 B2	4/2005 Brewer et al.	7,984,259 B1	7/2011	English
	6,880,062 B1	4/2005 Ibrahim et al.	8,031,703 B2	10/2011	Gottumukkula et al.
	6,898,670 B2	5/2005 Nahum	8,032,621 B1	10/2011	Upalekar et al.
			8,051,197 B2	11/2011	Mullendore et al.
			8,086,755 B2	12/2011	Duffy, IV et al.
			8,161,134 B2	4/2012	Mishra et al.
			8,196,018 B2	6/2012	Forhan et al.
			8,205,951 B2	6/2012	Boks
			8,218,538 B1	7/2012	Chidambaram et al.

(56)

References Cited

U.S. PATENT DOCUMENTS

8,230,066 B2	7/2012	Heil	9,448,877 B2	9/2016	Candelaria
8,234,377 B2	7/2012	Cohn	9,471,348 B2	10/2016	Zuo et al.
8,266,238 B2	9/2012	Zimmer et al.	9,501,473 B1	11/2016	Kong et al.
8,272,104 B2	9/2012	Chen et al.	9,503,523 B2	11/2016	Rosset et al.
8,274,993 B2	9/2012	Sharma et al.	9,565,110 B2	2/2017	Mullendore et al.
8,290,919 B1	10/2012	Kelly et al.	9,575,828 B2	2/2017	Agarwal et al.
8,297,722 B2	10/2012	Chambers et al.	9,582,377 B1	2/2017	Dhoolam et al.
8,301,746 B2	10/2012	Head et al.	9,614,763 B2	4/2017	Dong et al.
8,335,231 B2	12/2012	Kloth et al.	9,658,868 B2	5/2017	Hill
8,341,121 B1	12/2012	Claudatos et al.	9,658,876 B2	5/2017	Chang et al.
8,345,692 B2	1/2013	Smith	9,686,209 B1	6/2017	Arad et al.
8,352,941 B1	1/2013	Protopopov et al.	9,733,868 B2	8/2017	Chandrasekaran et al.
8,392,760 B2	3/2013	Kandula et al.	9,763,518 B2	9/2017	Charest et al.
8,442,059 B1	5/2013	De La Iglesia et al.	9,830,240 B2	11/2017	George et al.
8,479,211 B1	7/2013	Marshall et al.	9,838,341 B1 *	12/2017	Kadosh H04L 49/9084
8,495,356 B2	7/2013	Ashok et al.	9,853,873 B2	12/2017	Dasu et al.
8,514,868 B2	8/2013	Hill	2002/0049980 A1	4/2002	Hoang
8,532,108 B2	9/2013	Li et al.	2002/0053009 A1	5/2002	Selkirk et al.
8,560,663 B2	10/2013	Baucke et al.	2002/0073276 A1	6/2002	Howard et al.
8,619,599 B1	12/2013	Even	2002/0083120 A1	6/2002	Soltis
8,626,891 B2	1/2014	Guru et al.	2002/0095547 A1	7/2002	Watanabe et al.
8,630,983 B2	1/2014	Sengupta et al.	2002/0103889 A1	8/2002	Markson et al.
8,660,129 B1	2/2014	Brendel et al.	2002/0103943 A1	8/2002	Lo et al.
8,661,299 B1	2/2014	Ip	2002/0112113 A1	8/2002	Karpoff et al.
8,677,485 B2	3/2014	Sharma et al.	2002/0120741 A1	8/2002	Webb et al.
8,683,296 B2	3/2014	Anderson et al.	2002/0138675 A1	9/2002	Mann
8,706,772 B2	4/2014	Hartig et al.	2002/0156971 A1	10/2002	Jones et al.
8,719,804 B2	5/2014	Jain	2003/0023885 A1	1/2003	Potter et al.
8,725,854 B2	5/2014	Edsall et al.	2003/0026267 A1	2/2003	Oberman et al.
8,768,981 B1	7/2014	Milne et al.	2003/0055933 A1	3/2003	Ishizaki et al.
8,775,773 B2	7/2014	Acharya et al.	2003/0056126 A1	3/2003	O'Connor et al.
8,793,372 B2	7/2014	Ashok et al.	2003/0065986 A1	4/2003	Fraenkel et al.
8,805,918 B1	8/2014	Chandrasekaran et al.	2003/0084359 A1	5/2003	Bresniker et al.
8,805,951 B1	8/2014	Faibish et al.	2003/0112814 A1 *	6/2003	Modali H04L 47/10 370/412
8,832,330 B1	9/2014	Lancaster	2003/0118053 A1	6/2003	Edsall et al.
8,855,116 B2	10/2014	Rosset et al.	2003/0131105 A1	7/2003	Czeiger et al.
8,856,339 B2	10/2014	Mestery et al.	2003/0131165 A1	7/2003	Asano et al.
8,868,474 B2	10/2014	Leung et al.	2003/0131182 A1	7/2003	Kumar et al.
8,887,286 B2	11/2014	Dupont et al.	2003/0140210 A1	7/2003	Testardi
8,898,385 B2	11/2014	Jayaraman et al.	2003/0149763 A1	8/2003	Heitman et al.
8,909,928 B2	12/2014	Ahmad et al.	2003/0154271 A1	8/2003	Baldwin et al.
8,918,510 B2	12/2014	Gmach et al.	2003/0159058 A1	8/2003	Eguchi et al.
8,918,586 B1	12/2014	Todd et al.	2003/0174725 A1	9/2003	Shankar
8,924,720 B2	12/2014	Raghuram et al.	2004/0024961 A1	2/2004	Cochran et al.
8,930,747 B2	1/2015	Levijarvi et al.	2004/0030857 A1	2/2004	Krakirian et al.
8,935,500 B1	1/2015	Gulati et al.	2004/0039939 A1	2/2004	Cox et al.
8,949,677 B1	2/2015	Brundage et al.	2004/0054776 A1	3/2004	Klotz et al.
8,996,837 B1	3/2015	Bono et al.	2004/0059807 A1	3/2004	Klotz et al.
9,003,086 B1	4/2015	Schuller et al.	2004/0117438 A1	6/2004	Considine et al.
9,007,922 B1	4/2015	Mittal et al.	2004/0123029 A1	6/2004	Dalal et al.
9,009,427 B2	4/2015	Sharma et al.	2004/0128470 A1	7/2004	Hetzler et al.
9,009,704 B2	4/2015	McGrath et al.	2004/0128540 A1	7/2004	Roskind
9,075,638 B2	7/2015	Barnett et al.	2004/0153863 A1	8/2004	Klotz et al.
9,112,818 B1	8/2015	Arad et al.	2004/0190901 A1	9/2004	Fang
9,141,554 B1	9/2015	Candelaria	2004/0215749 A1	10/2004	Tsao
9,141,785 B2	9/2015	Mukkara et al.	2004/0230848 A1	11/2004	Mayo et al.
9,164,795 B1	10/2015	Mncent	2004/0250034 A1	12/2004	Yagawa et al.
9,176,677 B1	11/2015	Fradkin et al.	2005/0033936 A1	2/2005	Nakano et al.
9,201,704 B2	12/2015	Chang et al.	2005/0036499 A1	2/2005	Dutt et al.
9,203,784 B2	12/2015	Chang et al.	2005/0050211 A1	3/2005	Kaul et al.
9,207,882 B2	12/2015	Rosset et al.	2005/0050270 A1	3/2005	Horn et al.
9,207,929 B2	12/2015	Katsura	2005/0053073 A1	3/2005	Kloth et al.
9,213,612 B2	12/2015	Candelaria	2005/0055428 A1	3/2005	Teraï et al.
9,223,564 B2	12/2015	Munireddy et al.	2005/0060574 A1	3/2005	Klotz et al.
9,223,634 B2	12/2015	Chang et al.	2005/0071851 A1	3/2005	Opheim
9,244,761 B2	1/2016	Yekhanin et al.	2005/0076113 A1	4/2005	Klotz et al.
9,250,969 B2	2/2016	Lagar-Cavilla et al.	2005/0091426 A1	4/2005	Horn et al.
9,264,494 B2	2/2016	Factor et al.	2005/0114611 A1	5/2005	Durham et al.
9,270,754 B2	2/2016	Iyengar et al.	2005/0114615 A1	5/2005	Ogasawara et al.
9,280,487 B2	3/2016	Candelaria	2005/0117522 A1	6/2005	Basavaiah et al.
9,304,815 B1	4/2016	Vasanth et al.	2005/0117562 A1	6/2005	Wrenn
9,313,048 B2	4/2016	Chang et al.	2005/0138287 A1	6/2005	Ogasawara et al.
9,374,270 B2	6/2016	Nakil et al.	2005/0169188 A1	8/2005	Cometto et al.
9,378,060 B2	6/2016	Jansson et al.	2005/0185597 A1	8/2005	Le et al.
9,396,251 B1	7/2016	Boudreau et al.	2005/0198523 A1	9/2005	Shanbhag et al.
			2005/0235072 A1	10/2005	Smith et al.
			2005/0283658 A1	12/2005	Clark et al.
			2006/0015928 A1	1/2006	Setty et al.

(56)

References Cited

U.S. PATENT DOCUMENTS

2006/0034302 A1	2/2006	Peterson	2011/0299539 A1	12/2011	Rajagopal et al.
2006/0045021 A1	3/2006	Deragon et al.	2011/0307450 A1	12/2011	Hahn et al.
2006/0075191 A1	4/2006	Lolayekar et al.	2011/0313973 A1	12/2011	Srivasa et al.
2006/0098672 A1	5/2006	Schzudin et al.	2012/0023319 A1	1/2012	Chin et al.
2006/0117099 A1	6/2006	Mogul	2012/0030401 A1	2/2012	Cowan et al.
2006/0136684 A1	6/2006	Le et al.	2012/0054367 A1	3/2012	Ramakrishnan et al.
2006/0184287 A1	8/2006	Belady et al.	2012/0072578 A1	3/2012	Alam
2006/0198319 A1	9/2006	Schondelmayer et al.	2012/0072985 A1	3/2012	Davne et al.
2006/0215297 A1	9/2006	Kikuchi	2012/0075999 A1	3/2012	Ko et al.
2006/0230227 A1	10/2006	Ogasawara et al.	2012/0084445 A1	4/2012	Brock et al.
2006/0242332 A1	10/2006	Johnsen et al.	2012/0084782 A1	4/2012	Chou et al.
2007/0005297 A1	1/2007	Beresniewicz et al.	2012/0096134 A1	4/2012	Suit
2007/0067593 A1	3/2007	Satoyama et al.	2012/0130874 A1	5/2012	Mane et al.
2007/0079068 A1	4/2007	Draggon	2012/0131174 A1	5/2012	Ferris et al.
2007/0091903 A1	4/2007	Atkinson	2012/0134672 A1	5/2012	Banerjee
2007/0094465 A1	4/2007	Sharma et al.	2012/0144014 A1	6/2012	Natham et al.
2007/0101202 A1	5/2007	Garbow	2012/0159112 A1	6/2012	Tokusho et al.
2007/0121519 A1	5/2007	Cuni et al.	2012/0167094 A1	6/2012	Suit
2007/0136541 A1	6/2007	Herz et al.	2012/0173589 A1	7/2012	Kwon et al.
2007/0162969 A1	7/2007	Becker	2012/0177039 A1	7/2012	Berman
2007/0211640 A1	9/2007	Palacharla et al.	2012/0177041 A1	7/2012	Berman
2007/0214316 A1	9/2007	Kim	2012/0177042 A1	7/2012	Berman
2007/0250838 A1	10/2007	Belady et al.	2012/0177043 A1	7/2012	Berman
2007/0258380 A1	11/2007	Chamdani et al.	2012/0177044 A1	7/2012	Berman
2007/0263545 A1	11/2007	Foster et al.	2012/0177045 A1	7/2012	Berman
2007/0276884 A1	11/2007	Hara et al.	2012/0177370 A1	7/2012	Berman
2007/0283059 A1	12/2007	Ho et al.	2012/0179909 A1	7/2012	Sagi et al.
2008/0016412 A1	1/2008	White et al.	2012/0201138 A1	8/2012	Yu et al.
2008/0034149 A1	2/2008	Sheen	2012/0210041 A1	8/2012	Flynn et al.
2008/0052459 A1	2/2008	Chang et al.	2012/0254440 A1	10/2012	Wang
2008/0059698 A1	3/2008	Kabir et al.	2012/0257501 A1	10/2012	Kucharczyk
2008/0114933 A1	5/2008	Ogasawara et al.	2012/0265976 A1	10/2012	Spieris et al.
2008/0126509 A1	5/2008	Subramanian et al.	2012/0281706 A1	11/2012	Agarwal et al.
2008/0126734 A1	5/2008	Murase	2012/0297088 A1	11/2012	Wang et al.
2008/0168304 A1	7/2008	Flynn et al.	2012/0303618 A1	11/2012	Dutta et al.
2008/0201616 A1	8/2008	Ashmore	2012/0307641 A1	12/2012	Arumilli et al.
2008/0244184 A1	10/2008	Lewis et al.	2012/0311106 A1	12/2012	Morgan
2008/0256082 A1	10/2008	Davies et al.	2012/0311568 A1	12/2012	Jansen
2008/0267217 A1	10/2008	Colville et al.	2012/0320788 A1	12/2012	Venkataramanan et al.
2008/0294888 A1	11/2008	Ando et al.	2012/0324114 A1	12/2012	Dutta et al.
2009/0063766 A1	3/2009	Matsumura et al.	2012/0331119 A1	12/2012	Bose et al.
2009/0083484 A1	3/2009	Basham et al.	2013/0003737 A1	1/2013	Sinicrope
2009/0094380 A1	4/2009	Qiu et al.	2013/0013664 A1	1/2013	Baird et al.
2009/0094664 A1	4/2009	Butler et al.	2013/0028135 A1	1/2013	Berman
2009/0125694 A1	5/2009	Innan et al.	2013/0036212 A1	2/2013	Jibbe et al.
2009/0201926 A1	8/2009	Kagan et al.	2013/0036213 A1	2/2013	Hasan et al.
2009/0222733 A1	9/2009	Basham et al.	2013/0054888 A1	2/2013	Bhat et al.
2009/0240873 A1	9/2009	Yu et al.	2013/0061089 A1	3/2013	Valiyaparambil et al.
2009/0282471 A1	11/2009	Green et al.	2013/0080823 A1	3/2013	Roth et al.
2009/0323706 A1	12/2009	Germain et al.	2013/0086340 A1	4/2013	Fleming et al.
2010/0011365 A1	1/2010	Gerovac et al.	2013/0100858 A1	4/2013	Kamath et al.
2010/0030995 A1	2/2010	Wang et al.	2013/0111540 A1	5/2013	Sabin
2010/0046378 A1	2/2010	Knapp et al.	2013/0138816 A1	5/2013	Kuo et al.
2010/0083055 A1	4/2010	Ozonat	2013/0138836 A1	5/2013	Cohen et al.
2010/0174968 A1	7/2010	Charles et al.	2013/0139138 A1	5/2013	Kakos
2010/0318609 A1	12/2010	Lahiri et al.	2013/0144933 A1	6/2013	Hinni et al.
2010/0318837 A1	12/2010	Murphy et al.	2013/0152076 A1	6/2013	Patel
2011/0010394 A1	1/2011	Carew et al.	2013/0152175 A1	6/2013	Hromoko et al.
2011/0022691 A1	1/2011	Banerjee et al.	2013/0163426 A1	6/2013	Beliveau et al.
2011/0029824 A1	2/2011	Scholer et al.	2013/0163606 A1	6/2013	Bagepalli et al.
2011/0035494 A1	2/2011	Pandey	2013/0179941 A1	7/2013	McGloin et al.
2011/0087848 A1	4/2011	Trent	2013/0182712 A1	7/2013	Aguayo et al.
2011/0119556 A1	5/2011	De Buen	2013/0185433 A1	7/2013	Zhu et al.
2011/0142053 A1	6/2011	Van Der Merwe et al.	2013/0191106 A1	7/2013	Kephart et al.
2011/0161496 A1	6/2011	Nicklin	2013/0208888 A1	8/2013	Agrawal et al.
2011/0173303 A1	7/2011	Rider	2013/0212130 A1	8/2013	Rahnama
2011/0185117 A1*	7/2011	Beeston	2013/0223236 A1	8/2013	Dickey
		G06F 3/0652	2013/0238641 A1	9/2013	Mandelstein et al.
		711/111	2013/0266307 A1	10/2013	Garg et al.
2011/0228679 A1	9/2011	Varma et al.	2013/0268922 A1	10/2013	Tiwari et al.
2011/0231899 A1	9/2011	Pulier et al.	2013/0275470 A1	10/2013	Cao et al.
2011/0239039 A1	9/2011	Dieffenbach et al.	2013/0297655 A1	11/2013	Narasayya et al.
2011/0252274 A1	10/2011	Kawaguchi et al.	2013/0318134 A1	11/2013	Bolik et al.
2011/0255540 A1	10/2011	Mizrahi et al.	2013/0318288 A1	11/2013	Khan et al.
2011/0276584 A1	11/2011	Cotner et al.	2014/0006708 A1	1/2014	Huynh et al.
2011/0276675 A1	11/2011	Singh et al.	2014/0016493 A1	1/2014	Johnsson et al.
			2014/0019684 A1	1/2014	Wei et al.
			2014/0025770 A1	1/2014	Warfield et al.
			2014/0029441 A1	1/2014	Nydell

(56)

References Cited**U.S. PATENT DOCUMENTS**

2014/0029442 A1 1/2014 Wallman
 2014/0039683 A1 2/2014 Zimmermann
 2014/0040473 A1 2/2014 Ho et al.
 2014/0040883 A1 2/2014 Tompkins
 2014/0047201 A1 2/2014 Mehta
 2014/0053264 A1 2/2014 Dubrovsky et al.
 2014/0059187 A1 2/2014 Rosset et al.
 2014/0059266 A1 2/2014 Ben-Michael et al.
 2014/0064079 A1 3/2014 Kwan et al.
 2014/0086253 A1 3/2014 Yong
 2014/0089273 A1 3/2014 Borshack et al.
 2014/0095556 A1 4/2014 Lee et al.
 2014/0105009 A1 4/2014 Vos et al.
 2014/0105218 A1* 4/2014 Anand H04L 47/6255
 370/412
 2014/0108474 A1 4/2014 David et al.
 2014/0109071 A1 4/2014 Ding
 2014/0112122 A1 4/2014 Kapadia et al.
 2014/0123207 A1 5/2014 Agarwal et al.
 2014/0156557 A1 6/2014 Zeng et al.
 2014/0164666 A1 6/2014 Yang
 2014/0164866 A1 6/2014 Bolotov et al.
 2014/0172371 A1 6/2014 Zhu et al.
 2014/0173060 A1 6/2014 Jubran et al.
 2014/0173579 A1 6/2014 McDonald et al.
 2014/0189278 A1 7/2014 Peng
 2014/0198794 A1 7/2014 Mehta et al.
 2014/0211661 A1 7/2014 Gorkemli et al.
 2014/0215265 A1 7/2014 Mohanta et al.
 2014/0215590 A1 7/2014 Brand
 2014/0219086 A1 8/2014 Cantu' et al.
 2014/0222953 A1 8/2014 Karve et al.
 2014/0229790 A1 8/2014 Goss et al.
 2014/0244585 A1 8/2014 Sivasubramanian et al.
 2014/0244897 A1 8/2014 Goss et al.
 2014/0245435 A1 8/2014 Belenky
 2014/0269390 A1 9/2014 Ciodaru et al.
 2014/0281700 A1 9/2014 Nagesharao et al.
 2014/0297941 A1 10/2014 Rajani et al.
 2014/0307578 A1 10/2014 DeSanti
 2014/0317206 A1 10/2014 Lomelino et al.
 2014/0324862 A1 10/2014 Bingham et al.
 2014/0325208 A1 10/2014 Resch et al.
 2014/0331276 A1 11/2014 Frascadore et al.
 2014/0348166 A1 11/2014 Yang et al.
 2014/0355450 A1 12/2014 Bhikkaji et al.
 2014/0366155 A1 12/2014 Chang et al.
 2014/0376550 A1 12/2014 Khan et al.
 2015/0003450 A1 1/2015 Salam et al.
 2015/0003458 A1 1/2015 Li et al.
 2015/0003463 A1 1/2015 Li et al.
 2015/0010001 A1 1/2015 Duda et al.
 2015/0016461 A1 1/2015 Qiang
 2015/0030024 A1 1/2015 Venkataswami et al.
 2015/0046123 A1 2/2015 Kato
 2015/0063353 A1 3/2015 Kapadia et al.
 2015/0067001 A1 3/2015 Koltsidas
 2015/0082432 A1 3/2015 Eaton et al.
 2015/0092824 A1 4/2015 Wicker, Jr. et al.
 2015/0120907 A1 4/2015 Niestemski et al.
 2015/0121131 A1 4/2015 Kiselev et al.
 2015/0127979 A1 5/2015 Doppalapudi
 2015/0142840 A1 5/2015 Baldwin et al.
 2015/0180672 A1 6/2015 Kuwata
 2015/0180787 A1 6/2015 Hazelet et al.
 2015/0205974 A1 7/2015 Talley et al.
 2015/0207763 A1 7/2015 Bertran Ortiz et al.
 2015/0222444 A1 8/2015 Sarkar
 2015/0229546 A1 8/2015 Somaiya et al.
 2015/0248366 A1 9/2015 Bergsten et al.
 2015/0248418 A1 9/2015 Bhardwaj et al.
 2015/0254003 A1 9/2015 Lee et al.
 2015/0254088 A1 9/2015 Chou et al.
 2015/0261446 A1 9/2015 Lee
 2015/0263993 A1 9/2015 Kuch et al.

2015/0269048 A1 9/2015 Marr et al.
 2015/0277804 A1 10/2015 Arnold et al.
 2015/0281067 A1 10/2015 Wu
 2015/0303949 A1 10/2015 Jafarkhani et al.
 2015/0341237 A1 11/2015 Cuni et al.
 2015/0341239 A1 11/2015 Bertran Ortiz et al.
 2015/0358136 A1 12/2015 Medard et al.
 2015/0379150 A1 12/2015 Duda
 2016/0004611 A1 1/2016 Lakshman et al.
 2016/0011936 A1 1/2016 Luby
 2016/0011942 A1 1/2016 Golbourn et al.
 2016/0054922 A1 2/2016 Awasthi et al.
 2016/0062820 A1 3/2016 Jones et al.
 2016/0070652 A1 3/2016 Sundaraman et al.
 2016/0087885 A1 3/2016 Tripathi et al.
 2016/0088083 A1 3/2016 Bharadwaj et al.
 2016/0119159 A1 4/2016 Zhao et al.
 2016/0119421 A1 4/2016 Semke et al.
 2016/0139820 A1 5/2016 Fluman et al.
 2016/0149639 A1 5/2016 Pham et al.
 2016/0205189 A1 7/2016 Mopur et al.
 2016/0210161 A1 7/2016 Rosset et al.
 2016/0231928 A1 8/2016 Lewis et al.
 2016/0274926 A1 9/2016 Narasimhamurthy et al.
 2016/0285760 A1 9/2016 Dong et al.
 2016/0292359 A1 10/2016 Tellis et al.
 2016/0294696 A1 10/2016 Gafni et al.
 2016/0294983 A1 10/2016 Kliteynik et al.
 2016/0334998 A1 11/2016 George et al.
 2016/0366094 A1 12/2016 Mason et al.
 2016/0378624 A1 12/2016 Jenkins, Jr. et al.
 2016/0380694 A1 12/2016 Guduru
 2017/0010874 A1 1/2017 Rosset et al.
 2017/0010930 A1 1/2017 Dutta et al.
 2017/0019475 A1 1/2017 Metz et al.
 2017/0068630 A1 3/2017 Iskandar et al.
 2017/0168970 A1 6/2017 Sajeeva et al.
 2017/0177860 A1 6/2017 Suarez et al.
 2017/0212858 A1 7/2017 Chu et al.
 2017/0273019 A1 9/2017 Park et al.
 2017/0277655 A1 9/2017 Das et al.
 2017/0337010 A1 11/2017 Kriss et al.
 2017/0337097 A1 11/2017 Sipos et al.
 2017/0340113 A1 11/2017 Charest et al.
 2017/0371558 A1 12/2017 George et al.
 2018/0097707 A1 4/2018 Wright et al.

FOREIGN PATENT DOCUMENTS

EP 2439637 A1 4/2012
 EP 2680155 A1 1/2014
 GB 2350028 B 5/2001
 WO 2004077214 A2 9/2004
 WO 2016003489 A1 1/2016

OTHER PUBLICATIONS

Aweya, James, et al., "Multi-level active queue management with dynamic thresholds," Elsevier, Computer Communications 25 (2002) pp. 756-771.
 Mahalingam M., et al., "Virtual extensible Local Area Network (VXLAN): A Framework for Overlaying Virtualized Layer 2 Networks over Layer 3 Networks," Independent Submission, RFC 7348, Aug. 2014, 22 pages, Retrieved on URL: <http://www.ietf.org/doc/rfc/rfc7348.html>.
 McQuerry S., "Cisco UCS M-Series Modular Servers for Cloud-Scale Workloads," White Paper, Cisco Systems Inc., Sep. 2014, 11 Pages.
 Monia C., et al., "IFCP—A Protocol for Internet Fibre Channel Networking: draft-monia-ips-ifcp-00.txt," Dec. 12, 2000, 6 Pages.
 Mueen A., et al., "Online Discovery and Maintenance of Time Series Motifs," KDD'10, The 16th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, Washington, DC, U.S.A., Jul. 25-28, 2010, 10 Pages.
 Muglia B., "Decoding SDN," Juniper Networks, Jan. 14, 2013, pp. 1-7, [Retrieved on Aug. 27, 2015] Retrieved from URL: <http://>

(56)

References Cited**OTHER PUBLICATIONS**

- forums.juniper.net/15/The-New-Network/Decoding-SDN/ba-p/174651.
- Murray J.F., et al., "Machine Learning Methods for Predicting Failures in Hard Drives: A Multiple-Instance Application," *Journal of Machine Learning Research* 6, May 2005, 34 Pages.
- Nelson M., "File Verification Using CRC," *Dr. Dobbs's Journal*, May 1, 1992, pp. 1-18, XP055130883.
- Pace A., "Technologies for Large Data Management in Scientific Computing," *International Journal of Modern Physics C*, Feb. 2014, vol. 25, No. 2, 72 Pages.
- Petersen C., "Introducing Lightning: A flexible NVMe JBOF," Mar. 9, 2016, 6 pages.
- Pinheiro E., et al., "Failure Trends in a Large Disk Drive Population," *FAST '07, 5th USENIX Conference on File and Storage Technologies*, San Jose, California, U.S.A., Feb. 13-16, 2007, 13 Pages.
- Raginsky M., et al., "Sequential Anomaly Detection in the Presence of Noise and Limited Feedback," *IEEE*, Mar. 13, 2012, 19 Pages, arXiv:0911.2904v4 [cs.LG].
- Saidi Ali G., et al., "Performance Validation of Network-Intensive Workloads on a Full-System Simulator," *Interaction between Operating System and Computer Architecture Workshop, (IOSCA2005)*, Austin, Texas, Oct. 2005, 10 pages.
- Sajassi A., et al., "A Network Virtualization Overlay Solution Using EVPN," *L2VPN Workgroup*, Nov. 10, 2014, 24 pages, Retrieved from URL: <http://tools.ietf.org/pdf/draft-ietf-bess-evpn-overlay-00.pdf>.
- Sajassi A., et al., "BGP MPLS Based Ethernet VPN," *Network Working Group*, Oct. 18, 2014, 52 Pages.
- Sajassi A., et al., "Integrated Routing and Bridging in EVPN", *L2VPN Workgroup*, Nov. 11, 2014, 26 Pages.
- Schroeder B., et al., "Disk Failures in the Real World: What Does an MTTF of 1,000,000 Hours Mean to You?," *USENIX Association, FAST '07: 5th USENIX Conference on File and Storage Technologies*, San Jose, California, U.S.A., Feb. 13-16, 2007, 16 Pages.
- Shue D., et al., "Performance Isolation and Fairness for Multi-Tenant Cloud Storage," *Usenix Association, 10th Usenix Symposium on Operating Systems Design Implementation*, 2012, 14 pages, Retrieved from URL: <https://www.usenix.org/system/files/conference/osdi12/osdi12-final-215.pdf>.
- Shunra, "Shunra for HP Software, Enabling Confidence in Application Performance Before Deployment," 2010, 2 pages.
- Staimer M., "Inside Cisco Systems' Unified Computing System," *Cisco*, Jul. 2009, 5 Pages, Retrieved from URL: <http://searchstorage.techtarget.com/report/Inside-Cisco-Systems-Unified-Computing-System>.
- Stamey J., et al., "Client-Side Dynamic Metadata in Web 2.0," *SIGDOC '07*, Oct. 22-24, 2007, pp. 155-161.
- Swami V., "Simplifying SAN Management for VMWare Boot from SAN, Utilizing Cisco UCS and Palo," Posted May 31, 2011, 6 Pages, Retrieved from URL: <http://virtualeverything.wordpress.com/2011/05/31/simplifying-san-management-for-vmware-boot-from-san-utilizing-cisco-ucs-and-palo/>.
- Tate J., et al., "Introduction to Storage Area Networks," *IBM, Redbooks*, Dec. 2017, vol. 9, 302 Pages, Retrieved from URL: ibm.com/redbooks.
- "VBlock Solution for SAP: Simplified Provisioning for Operation Efficiency," *VCE White Paper, VCE—The Virtual Computing Environment Company*, Aug. 2011, 11 Pages, © 2011 VCE Company LLC, All Rights reserved, Retrieved from URL: <http://www.vce.com/pdf/solutions/vce-sap-simplified-provisioning-white-paper.pdf>.
- Vuppala V., et al., "Layer-3 Switching Using Virtual Network Ports," *Computer Communications and Networks, Proceedings of Eighth International Conference in Boston, MA, USA, Piscataway, NJ, USA, IEEE*, Oct. 11-13, 1999, pp. 642-648, ISBN: 0-7803-5794-9.
- Wang F., et al., "OBFS: A File System for Object-Based Storage Devices," *Storage System Research Center, MSST*, Apr. 2004, vol. 4, 18 Pages.
- Weil S.A., "Ceph: Reliable, Scalable, and High-Performance Distributed Storage," *University of California, Santa Cruz, on the Cisco Catalyst 4500 Classic Supervisor Engines*, Dec. 2007, 239 Pages, Retrieved from the Internet: URL: <https://ceph.com/papers/weil-thesis.pdf>.
- Weil S.A., et al., "Ceph: A Scalable, High-performance Distributed File System," *Proceedings of the 7th symposium on Operating systems design and implementation, USENIX Association*, Nov. 6, 2006, 14 Pages.
- Weil S.A., et al., "CRUSH: Controlled, Scalable, Decentralized Placement of Replicated Data," *Proceedings of the 2006 ACM/IEEE conference on Supercomputing, ACM*, Nov. 11, 2006, 12 Pages.
- Wikipedia, the Free Encyclopedia: "Standard RAID Levels," Last updated Jul. 18, 2014, 7 Pages, Retrieved from URL: http://en.wikipedia.org/wiki/Standard_RAID_levels.
- Wu J., et al., "The Design, and Implementation of AQUA: An Adaptive Object-Based Storage Device," *Department of Computer Science, MSST*, May 17, 2006, 25 Pages, Retrieved from URL: <http://storageconference.us/2006/Presentations/30Wu.pdf>.
- Xue C., et al., "A Standard Framework for Ceph Performance Profiling With Latency Breakdown," *CEPH*, Jun. 30, 2015, 3 Pages.
- Zhou Z., et al., "Stable Principal Component Pursuit," Jan. 14, 2010, 5 Pages, arXiv:1001.2363v1 [cs.IT].
- Zhu Y., et al., "A Cost-based Heterogeneous Recovery Scheme for Distributed Storage Systems with RAID-6 Codes," *University of Science & Technology of China*, 2012, 12 Pages.
- "Appendix D: Configuring In-Band Management," *Sun Storage Common Array Manager Installation and Setup Guide, Version 6.7.x 821-1362-10*, Sun Oracle, Copyright © 2010, Oracle and/or its affiliates, 15 Pages, [Retrieved and Printed Sep. 12, 2013], Retrieved from URL: <http://docs.oracle.com/cd/E19377-01/821-1362-10/index.html>.
- Author Unknown, "5 Benefits of a Storage Gateway in the Cloud," *Blog, TwinStrata, Inc.*, Jul. 25, 2012, XP055141645, 4 pages.
- Author Unknown, "Coraid EtherCloud, Software-Defined Storage with Scale-Out Infrastructure," *Solution Brief, Coraid, Redwood City, California, U.S.A.*, 2013, 2 pages.
- Author Unknown, "Creating Performance-Based SAN SLAs Using Finisar's NetWisdom," *Finisar Corporation, Sunnyvale, California, U.S.A.*, May 2006, 7 pages.
- Author Unknown, "Data Center, Metro Cloud Connectivity: Integrated Metro SAN Connectivity in 16 Gbps Switches," *Brocade Communication Systems, Inc.*, Apr. 2011, 14 pages.
- Author Unknown, "Data Center, San Fabric Administration Best Practices Guide, Support Perspective," *Brocade Communication Systems, Inc.*, May 2013, 21 pages.
- Author Unknown, "Delphi-Save a CRC Value in a File, Without Altering the Actual CRC Checksum?," *Stack Overflow, Stackoverflow.com*, Dec. 23, 2011, 3 pages, XP055130879, Retrieved from URL: <http://stackoverflow.com/questions/8608219/save-a-crc-value-in-a-file-without-altering-the-actual-crc-checksum>.
- Author Unknown, "EMC Unisphere: Innovative Approach to Managing Low-End and Midrange Storage, Redefining Simplicity in the Entry-Level and Midrange Storage Markets," *Data Sheet, EMC Corporation*, Published on or about Jan. 4, 2013, 6 pages, [Retrieved on Sep. 12, 2013], Retrieved from URL: <http://www.emc.com/storage/vnx/unisphere.html>.
- Author Unknown, "Joint Cisco and VMWare Solution for Optimizing Virtual Desktop Delivery: Data Center 3.0: Solutions to Accelerate Data Center Virtualization," *Cisco Systems Incorporated, VMware Incorporated*, Sep. 2008, 10 Pages.
- Author Unknown, "Network Transformation with Software-Defined Networking and Ethernet Fabrics," *Positioning Paper*, 2012, 6 pages, *Brocade Communications Systems*.
- Author Unknown, "Software Defined Networking: The New Norm for Networks," *ONF White Paper, Open Networking Foundation*, Apr. 13, 2012, 12 Pages.

(56)

References Cited**OTHER PUBLICATIONS**

Author Unknown, "Software Defined Storage Networks an Introduction," White Paper, Doc # 01-000030-001 Rev. A, Jeda Networks, Newport Beach, California, U.S.A., Dec. 12, 2012, 8 pages.

Author Unknown, "The Fundamentals of Software-Defined Storage, Simplicity at Scale for Cloud-Architectures," Solution Brief, Coraid, Redwood City, California, U.S.A., 2013, 3 pages.

Author Unknown, "Storage Area Network—NPIV: Emulex Virtual HBA and Brocade, Proven Interoperability and Proven Solution," Emulex and Brocade Communications Systems, Technical Brief, Apr. 2008, 4 Pages.

Author Unknown, "Storage Infrastructure for the Cloud," Solution Brief, Coraid, Redwood City, California, U.S.A., 2012, 3 pages.

Author Unknown, "Virtual Wisdom SAN Performance Probe Family Models: Probe FC8, Hd, and HD48," Virtual Instruments DataSheet, Apr. 2014, 4 Pages.

Author Unknown, "Xgig Analyzer: Quick Start Feature Guide 4.0," Finisar Corporation, Sunnyvale, California, U.S.A., Feb. 2008, 24 Pages.

Berman S., et al., "Start-Up Jeda Networks in Software Defined Storage Networks Technology," Storage News Letter, Press Release, Feb. 25, 2013, 2 pages, [Retrieved from Jun. 16, 2013] Retrieved from URL: <http://storagenewsletter.com/news/startups/jeda-networks>.

Borovick L., et al., "Architecting the Network for the Cloud," IDC White Paper, Jan. 2011, 8 pages.

Chakrabarti K., et al., "Locally Adaptive Dimensionality Reduction for Indexing Large Time Series Databases," ACM Transactions on Database Systems, Jun. 2002, vol. 27, No. 2, pp. 188-228.

Chandola V., et al., "A Gaussian Process Based Online Change Detection Algorithm for Monitoring Periodic Time Series," Proceedings of the Eleventh SIAM International Conference on Data Mining, SDM, Apr. 28-30, 2011, 12 Pages.

Choudry A.K., et al., "Dynamic Queue Length Thresholds for Shared-Memory Packet Switches," IEEE/ACM Transactions on Networking, Apr. 1998, vol. 6, No. 2, 29 Pages.

Cisco "N-Port Virtualization in the Data Center," Cisco White Paper, Cisco Systems, Inc., Mar. 2008, 6 Pages, Retrieved from URL: <http://www.cisco.com/en/US/prod/collateral/ps4159/ps6409/ps5989/ps9898/white-paper-c11-459263.html>.

Cisco Systems Inc., "Best Practices in Deploying Cisco Nexus 1000V Series Switches on Cisco UCS B and C Series Cisco UCS Manager Servers," Cisco White Paper, Apr. 2011, Retrieved from URL: <http://www.cisco.com/en/US/prod/collateral/switches/ps9441/ps9902/whitepaperc11-558242.pdf>, 36 Pages.

Cisco Systems Inc: "Cisco Prime Data Center Network Manager 6.1," At-a-Glance, 2012, 3 Pages.

Cisco Systems Inc: "Cisco Prime Data Center Network Manager: Release 6.1," Data Sheet, 2012, 10 Pages.

Cisco Systems Inc., "Cisco Unified Network Services: Overcome Obstacles to Cloud-Ready Deployments," Cisco White Paper, Jan. 2011, 6 Pages.

Clarke A., et al., "Open Data Center Alliance Usage: Virtual Machine (VM) Interoperability in a Hybrid Cloud Environment Rev. 1.2," Open Data Center Alliance, 2013, pp. 1-18.

"Configuration Interface for IBM System Storage DS5000, IBM DS4000, and IBM DS3000 Systems," IBM SAN Volume Controller Version 7.1, IBM® System Storage® SAN Volume Controller Information Center, Jun. 16, 2013, 3 Pages, Retrieved from URL: <http://publib.boulder.ibm.com/infocenter/svc/ic/index.jsp?topic=%2Fcom.ibm.storage.svc.console.doc%2Fsvc-fastcontconfint-1ev5ej.html>.

"Coraid Virtual Das (VDAS) Technology, Eliminate Tradeoffs Between DAS and Networked Storage," Coraid Technology Brief, 2013, © Coraid, Inc., Published on or about Mar. 20, 2013, 2 pages, Retrieved from URL: <http://san.coraid.com/rs/coraid/images/TechBrief-Coraid-VDAS.pdf>.

Cummings R., et al., "Fibre Channel- Fabric Generic Requirements (FC-FG)," American National Standards Institute, Inc., New York, U.S.A., Dec. 4, 1996, 33 Pages.

Farber F., et al., "An In-Memory Database System for Multi-Tenant Applications," Proceedings of 14th Business, Technology and Web (BTW) Conference on "Database Systems for Business, Technology, and Web", University of Kaiserslautern, Germany, Feb. 28-Mar. 4, 2011, 17 Pages, Retrieved from URL: <http://cs.emis.de/LNI/Proceedings/Proceedings180/650.pdf>.

Guo C.J., et al., "IBM Research Report: Data Integration and Composite Business Services, Part 3, Building a Multi-Tenant Data Tier with with Access Control and Security," RC24426 (C0711-037), Computer Science, Nov. 19, 2007, 20 Pages.

Hatzieleftheriou A., et al., "Host-side Filesystem Journaling for Durable Shared Storage," USENIX Association, 13th USENIX Conference on File and Storage Technologies (FAST '15), Feb. 16-19, 2015, 9 Pages, Retrieved from URL: <https://www.usenix.org/system/files/conference/fast15/fast15-paper-hatzieleftheriou.pdf>.

Hedayat K., et al., "A Two-Way Active Measurement Protocol (TWAMP)," Network Working Group, RFC 5357, <https://www.rfc-editor.org/rfc/rfc5357.html>, Oct. 2008, 26 Pages.

Horn C., et al., "Online Anomaly Detection With Expert System Feedback in Social Networks," IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), Prague, [Abstract only], May 22-27, 2011, 2 pages.

Hosterman C., et al., "Using EMC Symmetrix Storage in VMware vSphere Environments," Version 8.0, EMC2Techbooks, EMC Corporation, Published on or about Jul. 8, 2008, 314 Pages, [Retrieved and Printed Sep. 12, 2013], Retrieved from URL: <http://www.emc.com/collateral/hardware/solution-overview/h2529-vmware-esx-svrtw-symmetrix-wp-ldv.pdf>.

"HP XP Array Manager Software-Overview Features," Storage Device Management Software, Hewlett-Packard Development Company, 2013, 3 pages, © Hewlett-Packard Development Company, L.P. [Retrieved on Sep. 12, 2013], Retrieved from URL: <http://h18006.www1.hp.com/products/storage/software/amxsp/index.html>.

Hu Y., et al., "Cooperative Recovery of Distributed Storage Systems from Multiple Losses with Network Coding," IEEE Journal on Selected Areas in Communications, University of Science & Technology of China, vol. 28, No. 2, Feb. 2010, 9 Pages.

International Preliminary Report on Patentability for International Application No. PCT/US2017/043463, mailed Mar. 14, 2019, 8 Pages.

Juniper Networks, Inc., "Recreating Real Application Traffic in Junosphere Lab," Solution Brief, Dec. 2011, 3 pages.

Keogh E., et al., "Dimensionality Reduction for Fast Similarity Search in Large Time Series Databases," Knowledge and Information Systems Long Paper, May 16, 2000, 19 Pages.

Kolyshkin K., "Virtualization in Linux," Sep. 1, 2006, 5 pages, XP055141648, Retrieved from URL: <https://web.archive.org/web/20070120205111/http://download.openvz.org/doc/openvz-intro.pdf>.

Kovar J.F., "Startup Jeda Networks Takes SDN Approach to Storage Networks," CRN Press Release, Feb. 22, 2013, 1 Page, [Retrieved on Jun. 16, 2013] Retrieved from URL: <http://www.crn.com/240149244/printablearticle.htm>.

Lampson B.W., et al., "Crash Recovery in a Distributed Data Storage System," Jun. 1, 1979, 28 Pages.

Lewis M.E., et al., "Design of an Advanced Development Model Optical Disk-Based Redundant Array of Independent Disks (RAID) High Speed Mass Storage Subsystem," Final Technical Report, Oct. 1997, pp. 1-211.

Lin J., et al., "Finding Motifs in Time Series," SIGKDD'02, Edmonton, Alberta, Canada, Jul. 23-26, 2002, 11 Pages.

Linthicum D., "VM Import Could be a Game Changer for Hybrid Clouds," InfoWorld, Dec. 23, 2010, 4 pages.

Long A., Jr., "Modeling the Reliability of RAID Sets," Dell Power Solutions, May 2008, 4 Pages.

Ma A., et al., "RAIDShield: Characterizing, Monitoring, and Proactively Protecting Against Disk Failures," FAST '15, 13th USENIX Conference on File and Storage Technologies, Santa Clara, California, U.S.A., Feb. 16-19, 2015, 17 Pages.

* cited by examiner




FIG. 1A

Time	Total Mem Occupancy	Total Free	Dynamic Queue Max	Qo occupancy	Q1 occupancy	Q2 occupancy
T0	0	100	100	0	0	0
T1	10	90	100	10	0	0
T2	90	10	90	90	0	0
T3	95	5	45	90	5	0
T4	98	2	18	90	8	0
T5	99	1	9	90	9	0
T6	100	0	0	90	9	1

FIG. 1B

Time	Total Mem Occupancy	Total Free	Dynamic Queue Max	Qo occupancy	Q1 occupancy
T0	0	100	100	0	0
T1	10	90	100	10	0
T2	67	33	81	67	0
T3	68	32	72	68	0
T4	69	31	63	69	0
T5	74	26	27	69	5
T6	76	24	9	69	7
T7	77	23	0	69	8

FIG. 1C




FIG. 2

301

Time	Total Mem Occupancy	Total Free	Dynamic Queue Max	Q0 occupancy	Q1 occupancy
T0	0	100	100	0	0
T1	10	90	100	10	0
T2	81	19	81	81	0
T3	86	14	36	81	5
T4	89	11	9	81	8
T5	90	10	0	81	9

FIG. 3A




FIG. 3B

FIG. 4

410




FIG. 5A





FIG. 5B



1

QUEUE PROTECTION USING A SHARED GLOBAL MEMORY RESERVE

CROSS-REFERENCE TO RELATED APPLICATIONS

This application is a continuation of U.S. patent application Ser. No. 15/250,860, filed on Aug. 29, 2016, of which is hereby expressly incorporated by reference in its entirety.

TECHNICAL FIELD

The subject technology pertains to managing memory resources in a network switch and in particular, for managing a shared buffer memory amongst multiple queues in a shared memory network switch.

BACKGROUND

Several different architectures are commonly used to build packet switches (e.g., IP routers, ATM switches and Ethernet switches). One architecture is the output queue (OQ) switch, which places received packets in various queues that are dedicated to outgoing ports. The packets are stored in their respective queues until it is their turn to depart (e.g. to be “popped”). While various types of OQ switches have different pros and cons, a shared memory architecture is one of the simplest techniques for building an OQ switch. In some implementations, a shared memory switch functions by storing packets that arrive at various input ports of the switch into a centralized shared buffer memory. When the time arrives for the packets to depart, they are read from the shared buffer memory and sent to an egress line.

There are various techniques for managing a shared memory buffer. In some memory management solutions, the network switch prevents any single output queue from taking more than a specified share of the buffer memory when the buffer is oversubscribed, and permits a single queue to take more than its share to handle incoming packet bursts if the buffer is undersubscribed.

BRIEF DESCRIPTION OF THE DRAWINGS

In order to describe the manner in which the above-recited and other advantages and features of the disclosure can be obtained, a more particular description of the principles briefly described above will be rendered by reference to specific embodiments thereof which are illustrated in the appended drawings. Understanding that these drawings depict only exemplary embodiments of the disclosure and are not therefore to be considered to be limiting of its scope, the principles herein are described and explained with additional specificity and detail through the use of the accompanying drawings in which:

FIG. 1A graphically illustrates an example of queue occupancy relative to free buffer memory in a shared memory network switch.

FIGS. 1B and 1C illustrate examples of memory allocation tables that indicate occupancy for various queues, as well as a total available free memory resource for a shared buffer memory.

FIG. 2 illustrates an example flow chart for implementing a shared buffer memory allocation algorithm utilizing a global shared reserve, according to some aspects of the technology.

2

FIG. 3A illustrates an example table of queue occupancy levels for multiple queues implementing a global shared reserve memory management technique, according to some aspects of the technology.

FIG. 3B graphically illustrates an example of memory allocated to a shared buffer memory by various queues using a global shared reserve memory management technique, according to some aspects of the technology.

FIG. 4 illustrates an example network device.

FIGS. 5A and 5B illustrate example system embodiments.

DESCRIPTION OF EXAMPLE EMBODIMENTS

Various embodiments of the disclosure are discussed in detail below. While specific implementations are discussed, it should be understood that this is done for illustration purposes only. A person skilled in the relevant art will recognize that other components and configurations can be used without parting from the spirit and scope of the disclosure.

Overview

One problem with managing shared memory space amongst multiple queues is to ensure that active queues (i.e., “aggressor queues”) do not occupy the entire memory and thereby impede buffer access by other queues. Queues that are prevented from enqueue due to limited buffer space are referred to herein as “victim queues.” In a shared memory switch, an algorithm is required to prevent any single queue from taking more than its fair allocation of shared memory. In some memory management solutions, the algorithm calculates a dynamic maximum threshold by multiplying the amount of unallocated/free memory in the shared memory by a parameter (e.g., “alpha”). Typically values of alpha range between 0.5 and 2.0.

With alpha set to 1.0 consider a single oversubscribed queue: the system stabilizes with the queue and the free memory both being the same size, i.e., the queue can consume only half of memory. With 2 oversubscribed queues the queues can each have $\frac{1}{3}^{rd}$ of the memory and $\frac{1}{3}^{rd}$ remains unallocated, and so on up to N oversubscribed queues, where each queue will have $1/(N+1)$ of the memory and $1/(N+1)$ will remain unallocated. An example of the relative memory allocation amongst multiple queues is discussed in further detail with respect to FIG. 1A, below.

In some data center deployments, the buffer is required to be able to absorb large data bursts into a single queue (e.g., incast burst absorption). So the “alpha” parameter (which is programmable), is usually set to greater than 1, typically 9 (e.g., 90% of the buffer). With this setting, few aggressor queues/concurrent bursts could consume the entire buffer, and any new incoming traffic is dropped (e.g. a tail-drop), affecting throughput for victim queues.

Another solution is to provide a dedicated memory allocation for each queue (e.g., a minimum reserve), and reduce the total shareable buffer space by the sum of all minimum reserves. Depending on implementation, this can result in carving out a section of the buffer memory that isn’t efficiently used. Additionally, the amount of reserved buffer space is a function of the number of ports and classes of service required, so as the number of ports/services scale, dedicated memory allocations become increasingly likely to deplete available memory.

DESCRIPTION

Aspects of the subject technology address the foregoing problem by providing memory management systems, meth-

3

ods and computer-executable instructions to facilitate packet storage using a shared buffer memory. In particular, the disclosed technology provides a packet enqueueing method which requires certain preconditions before a received packet can be enqueued. In some aspects, the decision of whether or not to enqueue a packet is first based on a fill level of the shared buffer memory. That is, if an occupancy of the queue in the shared buffer memory is below a pre-determined dynamic queue threshold (e.g., a “dynamic queue maximum” or “dynamic queue MAX”), then the packet is enqueued.

Alternatively, in instances where the queue occupancy in the shared buffer exceeds the dynamic queue max threshold, then further conditions may be verified before the packet is enqueued (or dropped). As discussed in further detail below, if the fill level of the queue in the shared buffer memory exceeds the dynamic queue max threshold, then an occupancy of the referring queue may be compared to static queue threshold (e.g., a “static queue minimum” or “static queue MIN”), to determine if the packet can still be enqueued.

As used herein, the dynamic queue maximum refers to a measure of shared buffer occupancy for the entire shared buffer memory. Thus, the dynamic queue max can be understood as a function of total free/available memory in the buffer. As discussed in further detail below, the static queue minimum threshold is a threshold that relates to a minimum amount of memory in the shared buffer that is allocated for use by victim queues.

FIG. 1A graphically illustrates an example of queue occupancy levels relative to a free shared buffer memory allocation in a network switch. For example, shared buffer 102 illustrates an example in which a shared buffer occupancy is maximally allocated at $\frac{1}{2}$ of the total memory capacity, i.e., for a single queue wherein $\alpha=1.0$. Shared buffer 104 illustrates an example of a total buffer allocation for two total queues, wherein the respective shared buffer memory allocation for each is $\frac{1}{3}^{rd}$ of the of the total buffer size. Similarly, shared buffer 106 illustrates an example of a shared buffer allocation amongst N total queues.

FIG. 1B illustrates an example memory allocation table 108 that indicates occupancy for multiple queues, as well as a total free memory for a shared buffer.

In particular, the example of FIG. 1B illustrates a memory management scheme in which any queue is permitted to occupy the entirety of shared buffer memory. As table 108 illustrates, this memory management method can be problematic due to the fact that aggressor queues can rapidly utilize the entirety of shared memory space, consequently halting the ability for victim queues to enqueue additional incoming packets.

By way of example, table 108 illustrates various occupancy levels for multiple queues (i.e., Q_0 , Q_1 , and Q_2), such that any individual queue is permitted to utilize all available free memory. This scenario is demonstrated, for example, at time=T5 where Q_0 , and Q_1 occupy 90% and 9% of the total memory, respectively (leaving a total free memory of 1%).

FIG. 1C illustrates an example memory allocation table 110 that indicates occupancy levels for multiple queues using a memory management technique that employs a dedicated “minimum reserve” for each respective queue. As discussed above, such solutions can also be sub-optimal due to the fact that some amount of shared buffer memory can be persistently reserved for inactive queues, even when memory resources are needed elsewhere. For example, using a per-queue memory reservation technique depicted by FIG.

4

1C (e.g., with $\alpha=9$) a total of 25% of the total buffer memory is reserved for various queues.

By way of example, table 110 illustrates this scenario at time=T7, where Q_0 occupancy is at 69 (e.g., 69% of the shared buffer size), and Q_1 occupancy is at 8 (e.g., 8% of the shared buffer size), however, dynamic queue max=0, indicating that free memory (e.g., total free=23) is no longer available to other aggressor queues. Therefore, in this scenario, a total of 23% of the shared buffer memory is unallocated if all victim queues are unutilized.

As discussed above, aspects of the subject technology address the foregoing limitations of conventional buffer memory management techniques, by providing a shared buffer memory in which packet enqueueing is dependent upon the verification of various conditions, for example, relating to a fill level of the shared buffer (e.g., a dynamic queue max threshold), as well as comparisons between a fill level of a referring queue and a threshold related to a reserved apportionment of buffer resources (e.g., a static queue min threshold).

FIG. 2 illustrates an example flow chart for implementing a method 200 for allocating shared buffer memory in a network switch. Method 200 begins at step 202, in which a determination is made as to whether there is any available (unallocated) memory in the shared buffer of a network switch. If it is determined that no free memory is available, method 200 proceeds to step 204, and any newly arriving packets are dropped.

Alternatively, if it is determined that the shared buffer memory contains unallocated space, method 200 proceeds to step 206, in which a determination is made as to whether any shared buffer space is available in the shared buffer memory.

If it is determined in step 206 that no memory in the shared buffer is available, then method 200 proceeds to step 208, in which a determination is made as to whether or not the occupancy of the referring queue is below a predetermined static queue minimum, e.g., a “static queue MIN” threshold, as discussed above. In some aspects, the static queue MIN threshold is a predetermined threshold used to define a minimum threshold, above which the received data/packets from a referring queue cannot be accepted into the shared buffer memory. As such, if in step 208 it is determined that the referring queue occupancy is not less than the static queue minimum, then method 200 proceeds to step 204 and incoming packet/s are dropped.

Alternatively, if in step 208 it is determined that the referring queue occupancy is less than the static queue MIN threshold, method 200 proceeds to step 212, and data from the referring queue is stored in a “reserved portion” of the shared buffer memory. It is understood herein that the reserved portion of buffer memory (or “global reserve”) refers to a logical allotment of memory space in the shared buffer. However, it is not necessary that the global reserve portions of memory be physically distinct memory spaces that are separate, for example, from various other regions in the shared memory buffer.

Referring back to step 206, if it is determined that shared memory space is available, then method 200 proceeds to step 210 in which a determination is made as to whether the queue occupancy is less than a dynamic queue threshold (e.g., “dynamic queue MAX”). As used herein, the dynamic queue max is a threshold that defines a cutoff, above which data from an aggressor queue cannot be admitted into the shared buffer memory. Because the dynamic queue max is a function of unallocated memory space in the shared buffer memory, in some aspects the dynamic queue max threshold

may be conceptualized as a function of queue activity for each associated queue in the network switch.

If in step 210 it is determined that the queue occupancy is less than the dynamic queue max, then method 200 proceeds to step 212 and the packet/s are stored in the buffer memory. Alternatively, if it is determined that the queue occupancy (e.g., queue allocation) is greater than the dynamic queue max threshold, then method 200 proceeds to step 208, where it is determined if the referring queue occupancy is less than the static queue minimum (see above).

By providing a global reserve buffer available to any referring queue that has less than a specified occupancy level, the subject memory management techniques permit data storage in the shared buffer by less active (victim) queues, even in instances where the majority of buffer storage space has been filed by aggressor queues.

FIG. 3A illustrates an example table 301 of queue occupancy levels for multiple queues implementing a shared memory management technique, as discussed above. By implementing the memory management scheme discussed with respect to method 200 above, victim queues with relatively low throughput (as compared to aggressor queues) can access portions of shared buffer memory that would otherwise be unavailable in other shared memory schemes. As illustrated in table 301, for example, at time=T5, victim queue (Q₁) is able to occupy some amount of the total buffer memory (e.g., 9%), although aggressor queue Q₁ has occupied most of the shared buffer. A similar example is graphically illustrated with respect to FIG. 3B.

Specifically, FIG. 3B illustrates an example of the apportionment of a memory 303 amongst multiple queues (e.g., first queue 309, second queue 311, and third queue 313), using a global shared reserve management technique, according to aspects of the subject technology. As illustrated, memory 303 is logically apportioned into a dynamic allocation 305 and a global reserve 307. As discussed above, dynamic allocation 305 can be a shared resource available to any aggressor queue until an occupancy level of that queue has reached a pre-determined threshold (e.g., a dynamic queue max). However, the global reserve 307 of memory 303 remains reserved for low volume or victim queues, so long as the occupancy of the referring queue does not exceed a pre-determined threshold governing storage to the global reserve (e.g., a static queue min threshold), as discussed above with respect to step 208 of method 200.

By way of example, the occupancy of buffer memory 303 in the example of FIG. 3B illustrates storage by three different queues. The storage of data pertaining to first queue 309 is managed solely within dynamic allocation 305. The storage of data pertaining to second queue 311 is shared amongst dynamic allocation 305 and global reserve 307, and data associated with third queue 313 is stored exclusively into global reserve 307.

As discussed above, storage of data from second queue 311 first began by storing data to dynamic allocation 305, until occupancy of dynamic allocation 305 was complete. After dynamic allocation 305 reached capacity, a determination was made as to whether the remaining data in second queue 311 was smaller than a static queue threshold, necessary to admit the data into the global reserve. Lastly, data from third queue 313, which could not have been stored to dynamic allocation 305 (due to its fill state), was exclusively stored into global reserve 307.

By maintaining global reserve 307 portion of buffer memory 303, the disclosed memory management technique provides for a minimal apportionment of shared buffer space that is continuously available to victim queues.

Example Devices

FIG. 4 illustrates an example network device 410 suitable for high availability and failover. Network device 410 includes a master central processing unit (CPU) 462, interfaces 468, and a bus 415 (e.g., a PCI bus). When acting under the control of appropriate software or firmware, the CPU 462 is responsible for executing packet management, error detection, and/or routing functions. The CPU 462 preferably accomplishes all these functions under the control of software including an operating system and any appropriate applications software. CPU 462 may include one or more processors 463 such as a processor from the Motorola family of microprocessors or the MIPS family of microprocessors. In an alternative embodiment, processor 463 is specially designed hardware for controlling the operations of router 410. In a specific embodiment, a memory 461 (such as non-volatile RAM and/or ROM) also forms part of CPU 462. However, there are many different ways in which memory could be coupled to the system.

The interfaces 468 are typically provided as interface cards (sometimes referred to as "line cards"). Generally, they control the sending and receiving of data packets over the network and sometimes support other peripherals used with the router 410. Among the interfaces that may be provided are Ethernet interfaces, frame relay interfaces, cable interfaces, DSL interfaces, token ring interfaces, and the like. In addition, various very high-speed interfaces may be provided such as fast token ring interfaces, wireless interfaces, Ethernet interfaces, Gigabit Ethernet interfaces, ATM interfaces, HSSI interfaces, POS interfaces, FDDI interfaces and the like. Generally, these interfaces may include ports appropriate for communication with the appropriate media. In some cases, they may also include an independent processor and, in some instances, volatile RAM. The independent processors may control such communications intensive tasks as packet switching, media control and management. By providing separate processors for the communications intensive tasks, these interfaces allow the master microprocessor 462 to efficiently perform routing computations, network diagnostics, security functions, etc.

Although the system shown in FIG. 4 is one specific network device of the present invention, it is by no means the only network device architecture on which the present invention can be implemented. For example, an architecture having a single processor that handles communications as well as routing computations, etc. is often used. Further, other types of interfaces and media could also be used with the router.

Regardless of the network device's configuration, it may employ one or more memories or memory modules (including memory 461) configured to store program instructions for the general-purpose network operations and mechanisms for roaming, route optimization and routing functions described herein. The program instructions may control the operation of an operating system and/or one or more applications, for example. The memory or memories may also be configured to store tables such as mobility binding, registration, and association tables, etc.

FIG. 5A and FIG. 5B illustrate example system embodiments. The more appropriate embodiment will be apparent to those of ordinary skill in the art when practicing the present technology. Persons of ordinary skill in the art will also readily appreciate that other system embodiments are possible.

FIG. 5A illustrates a conventional system bus computing system architecture 500 wherein the components of the

system are in electrical communication with each other using a bus 505. Exemplary system 500 includes a processing unit (CPU or processor) 510 and a system bus 505 that couples various system components including the system memory 515, such as read only memory (ROM) 520 and random access memory (RAM) 525, to the processor 510. The system 500 can include a cache of high-speed memory connected directly with, in close proximity to, or integrated as part of the processor 510. The system 500 can copy data from the memory 515 and/or the storage device 530 to the cache 512 for quick access by the processor 510. In this way, the cache can provide a performance boost that avoids processor 510 delays while waiting for data. These and other modules can control or be configured to control the processor 510 to perform various actions. Other system memory 515 may be available for use as well. The memory 515 can include multiple different types of memory with different performance characteristics. The processor 510 can include any general purpose processor and a hardware module or software module, such as module 1 532, module 2 534, and module 3 536 stored in storage device 530, configured to control the processor 510 as well as a special-purpose processor where software instructions are incorporated into the actual processor design. The processor 510 may essentially be a completely self-contained computing system, containing multiple cores or processors, a bus, memory controller, cache, etc. A multi-core processor may be symmetric or asymmetric.

To enable user interaction with the computing device 500, an input device 545 can represent any number of input mechanisms, such as a microphone for speech, a touch-sensitive screen for gesture or graphical input, keyboard, mouse, motion input, speech and so forth. An output device 535 can also be one or more of a number of output mechanisms known to those of skill in the art. In some instances, multimodal systems can enable a user to provide multiple types of input to communicate with the computing device 500. The communications interface 540 can generally govern and manage the user input and system output. There is no restriction on operating on any particular hardware arrangement and therefore the basic features here may easily be substituted for improved hardware or firmware arrangements as they are developed.

Storage device 530 is a non-volatile memory and can be a hard disk or other types of computer readable media which can store data that are accessible by a computer, such as magnetic cassettes, flash memory cards, solid state memory devices, digital versatile disks, cartridges, random access memories (RAMs) 525, read only memory (ROM) 520, and hybrids thereof.

The storage device 530 can include software modules 532, 534, 536 for controlling the processor 510. Other hardware or software modules are contemplated. The storage device 530 can be connected to the system bus 505. In one aspect, a hardware module that performs a particular function can include the software component stored in a computer-readable medium in connection with the necessary hardware components, such as the processor 510, bus 505, display 535, and so forth, to carry out the function.

FIG. 5B illustrates an example computer system 550 having a chipset architecture that can be used in executing the described method and generating and displaying a graphical user interface (GUI). Computer system 550 is an example of computer hardware, software, and firmware that can be used to implement the disclosed technology. System 550 can include a processor 555, representative of any number of physically and/or logically distinct resources

capable of executing software, firmware, and hardware configured to perform identified computations. Processor 555 can communicate with a chipset 560 that can control input to and output from processor 555. In this example, chipset 560 outputs information to output device 565, such as a display, and can read and write information to storage device 570, which can include magnetic media, and solid state media, for example. Chipset 560 can also read data from and write data to RAM 575. A bridge 580 for interfacing with a variety of user interface components 585 can be provided for interfacing with chipset 560. Such user interface components 585 can include a keyboard, a microphone, touch detection and processing circuitry, a pointing device, such as a mouse, and so on. In general, inputs to system 550 can come from any of a variety of sources, machine generated and/or human generated.

Chipset 560 can also interface with one or more communication interfaces 590 that can have different physical interfaces. Such communication interfaces can include interfaces for wired and wireless local area networks, for broadband wireless networks, as well as personal area networks. Some applications of the methods for generating, displaying, and using the GUI disclosed herein can include receiving ordered datasets over the physical interface or be generated by the machine itself by processor 555 analyzing data stored in storage 570 or 575. Further, the machine can receive inputs from a user via user interface components 585 and execute appropriate functions, such as browsing functions by interpreting these inputs using processor 555.

It can be appreciated that example systems 500 and 550 can have more than one processor 510 or be part of a group or cluster of computing devices networked together to provide greater processing capability.

For clarity of explanation, in some instances the present technology may be presented as including individual functional blocks including functional blocks comprising devices, device components, steps or routines in a method embodied in software, or combinations of hardware and software.

In some embodiments the computer-readable storage devices, mediums, and memories can include a cable or wireless signal containing a bit stream and the like. However, when mentioned, non-transitory computer-readable storage media expressly exclude media such as energy, carrier signals, electromagnetic waves, and signals per se.

Methods according to the above-described examples can be implemented using computer-executable instructions that are stored or otherwise available from computer readable media. Such instructions can comprise, for example, instructions and data which cause or otherwise configure a general purpose computer, special purpose computer, or special purpose processing device to perform a certain function or group of functions. Portions of computer resources used can be accessible over a network. The computer executable instructions may be, for example, binaries, intermediate format instructions such as assembly language, firmware, or source code. Examples of computer-readable media that may be used to store instructions, information used, and/or information created during methods according to described examples include magnetic or optical disks, flash memory, USB devices provided with non-volatile memory, networked storage devices, and so on.

Devices implementing methods according to these disclosures can comprise hardware, firmware and/or software, and can take any of a variety of form factors. Typical examples of such form factors include laptops, smart phones, small form factor personal computers, personal

digital assistants, rackmount devices, standalone devices, and so on. Functionality described herein also can be embodied in peripherals or add-in cards. Such functionality can also be implemented on a circuit board among different chips or different processes executing in a single device, by way of further example.

The instructions, media for conveying such instructions, computing resources for executing them, and other structures for supporting such computing resources are means for providing the functions described in these disclosures.

Although a variety of examples and other information was used to explain aspects within the scope of the appended claims, no limitation of the claims should be implied based on particular features or arrangements in such examples, as one of ordinary skill would be able to use these examples to derive a wide variety of implementations. Further and although some subject matter may have been described in language specific to examples of structural features and/or method steps, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to these described features or acts. For example, such functionality can be distributed differently or performed in components other than those identified herein. Rather, the described features and steps are disclosed as examples of components of systems and methods within the scope of the appended claims. Moreover, claim language reciting “at least one of” a set indicates that one member of the set or multiple members of the set satisfy the claim.

What is claimed is:

1. A method of managing memory in a network switch, the method comprising:
 - receiving a data packet at a first network queue from among a plurality of network queues to determine whether to enqueue or drop the data packet;
 - determining if a fill level of the first network queue in a shared buffer of the network switch exceeds a dynamic queue threshold, wherein the dynamic queue threshold is a maximum threshold;
 - in response to the fill level of the first network queue in the shared buffer exceeding the dynamic queue threshold, determining if the fill level of the first network queue is less than a static queue minimum threshold; and
 - in response, determining whether to enqueue or drop the data packet.
2. The method of claim 1, further comprising: enqueueing the data packet in the shared buffer when the fill level of the first network queue is less than the static queue minimum threshold.
3. The method of claim 1, further comprising: dropping the data packet when the fill level of the first network queue is greater than the static queue minimum threshold.
4. The method of claim 1, further comprising: dropping the data packet when the fill level of the first network queue exceeds the dynamic queue threshold and the static queue minimum threshold.
5. The method of claim 1, wherein the dynamic queue threshold is based on an amount of unallocated memory in the shared buffer.
6. The method of claim 1, wherein the dynamic queue threshold is a function of a fill level for each respective one of the plurality of network queues.
7. The method of claim 1, wherein the static queue minimum threshold is a user defined value.
8. A shared memory network switch comprising: at least one processor;

a shared buffer memory, the shared buffer memory comprising a dynamic memory allocation and a reserve memory allocation;

and a memory device storing instructions that, when executed by the at least one processor, cause the at least one processor to perform operations comprising:

receiving a data packet at a first network queue from among a plurality of network queues to determine whether to enqueue or drop the data packet;

determining if a fill level of the first network queue in the shared buffer of the network switch exceeds a dynamic queue threshold, wherein the dynamic queue threshold is a maximum threshold;

and when the fill level of the first network queue in the shared buffer exceeds the dynamic queue threshold, determining if the fill level of the first network queue is less than a static queue minimum threshold; and in response, determining whether to enqueue or drop the data packet.

9. The shared memory network switch of claim 8, further comprising:

enqueueing the data packet in the shared buffer when the fill level of the first network queue is less than the static queue minimum threshold.

10. The shared memory network switch of claim 8, further comprising: dropping the data packet when the fill level of the first network queue is greater than the static queue minimum threshold.

11. The shared memory network switch of claim 8, further comprising:

dropping the data packet when the fill level of the first network queue exceeds the dynamic queue threshold and the static queue minimum threshold.

12. The shared memory network switch of claim 8, wherein the dynamic queue threshold is based on an amount of unallocated memory in the shared buffer.

13. The shared memory network switch of claim 8, wherein the dynamic queue threshold is a function of a fill level for each respective one of the plurality of network queues.

14. The shared memory network switch of claim 8, wherein the static queue minimum threshold is a user defined value.

15. A non-transitory computer-readable storage medium comprising instructions stored therein, which when executed by one or more processors, cause the processors to perform operations comprising:

receiving a data packet at a first network queue from among a plurality of network queues to determine whether to enqueue or drop the data packet;

determining if a fill level of the first network queue in a shared buffer exceeds a dynamic queue threshold, wherein the dynamic queue threshold is a maximum threshold;

and when the fill level of the first network queue in the shared buffer exceeds the dynamic queue threshold, determining if the fill level of the first network queue is less than a static queue minimum threshold; and in response, determining whether to enqueue or drop the data packet.

16. The non-transitory computer-readable storage medium of claim 15, further comprising:

enqueueing the data packet in the shared buffer when the fill level of the first network queue is less than the static queue minimum threshold.

17. The non-transitory computer-readable storage medium of claim 15, further comprising:

11

dropping the data packet when the fill level of the first network queue is greater than the static queue minimum threshold.

18. The non-transitory computer-readable storage medium of claim 15, further comprising: 5

dropping the data packet when the fill level of the fill level of the first network queue exceeds the dynamic queue threshold and the static queue minimum threshold.

19. The non-transitory computer-readable storage medium of claim 15, wherein the dynamic queue threshold 10 is based on an amount of unallocated memory in the shared buffer.

20. The non-transitory computer-readable storage medium of claim 15, wherein the dynamic queue threshold is a function of a fill level for each respective one of the 15 plurality of network queues.

* * * * *

12