

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12413587
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	Gross; David Steven et al.

Application identification

Abstract

This disclosure describes techniques for identifying an application (e.g., accessing application) that is attempting to access a resource. In some examples, access may be managed by an authentication service. When an access request is received at the authentication service from an application on a client device, the authentication service may ask the application to communicate with an identification agent on the client device. The identification agent may perform one or more tests to discover the identity of the application. In some cases, the identification agent may send the identity of the application to the authentication service. The authentication service may then allow or deny access by the accessing application to the resource based at least in part on the discovered identity.

Inventors: Gross; David Steven (Cincinnati, OH), Bammel; Jennifer Lee (Dexter, MI), Matteson; David William (Dexter, MI), Cassell; Christopher Carl (Ann Arbor, MI), Mills; Kyle David (Ann Arbor, MI)

Applicant: Cisco Technology, Inc. (San Jose, CA)

Family ID: 83117587

Assignee: Cisco Technology, Inc. (San Jose, CA)

Appl. No.: 18/477063

Filed: September 28, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240031365 A1	Jan. 25, 2024

Related U.S. Application Data

division parent-doc US 17194000 20210305 US 11799856 child-doc US 18477063

Publication Classification

Int. Cl.: H04L9/40 (20220101); H04L101/663 (20220101)

U.S. Cl.:

CPC H04L63/0876 (20130101); H04L2101/663 (20220501)

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6182146	12/2000	Graham-Cumming, Jr.	709/224	H04L 69/08
2013/0031254	12/2012	Bhandiwad	709/225	H04L 47/34
2014/0075513	12/2013	Trammel	726/4	G06F 21/41
2015/0020220	12/2014	Mullick et al.	N/A	N/A
2018/0218145	12/2017	Hussain	N/A	H04L 63/10
2018/0337920	12/2017	Stites	N/A	G06F 21/73
2018/0351986	12/2017	Johns	N/A	N/A
2020/0374280	12/2019	Zeng et al.	N/A	N/A
2022/0286455	12/2021	Gross et al.	N/A	N/A

OTHER PUBLICATIONS

Browser Detection Using the User Agent, Last Modified Dec. 23, 2020 and Downloaded Jan. 11, 2021, from MDN Web Docs Mozilla, at <<<https://developer.mozilla.org/en-US/docs/Web/HTTP/Browser_detection_using_the_user_agent>>>. cited by applicant
Office Action for U.S. Appl. No. 17/194,000, mailed on Feb. 16, 2023, David Steven Gross, "Application Identification", 6 pages. cited by applicant
Veerasundar, "How to check which application is using which port", retrieved from <<<https://veerasundar.com/blog/how-to-check-which-application-is-using-which-port/>>> on Feb. 10, 2023, 3 pages. cited by applicant

Primary Examiner: Christensen; Scott B

Attorney, Agent or Firm: Lee & Hayes, P.C.

Background/Summary

RELATED APPLICATIONS (1) This application is a divisional of and claims priority to U.S. patent application Ser. No. 17/194,000, filed Mar. 5, 2021.

TECHNICAL FIELD

(1) The present disclosure relates generally to identifying an application that is attempting to access a resource and/or service over a network, thereby improving performance and/or security of the

network.

BACKGROUND

(2) A user of a client device may wish to gain access to a resource and/or service (e.g., website) via a network. A decision to grant or deny access to the client device may be based on a policy. In some examples, the policy may be based at least in part on the identity of an application that is being used by the client device to access the network. However, if the application is not correctly identified, access may be granted inappropriately, or access may be denied inappropriately. In some examples, identification of the accessing application may be difficult. For example, a user agent on the client device may incorrectly report an identity of the accessing application, or may incorrectly report other information related to the application or related to the client device. Incorrect reporting of information about the application or client device may be accidental or may be purposeful, even nefarious. Regardless, improved methods for correctly identifying an accessing application may help make more informed decisions about whether to grant access to a resource.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The detailed description is set forth below with reference to the accompanying figures. In the figures, the left-most digit(s) of a reference number identifies the figure in which the reference number first appears. The use of the same reference numbers in different figures indicates similar or identical items. In some cases, parentheticals are utilized after a reference number to distinguish like elements. Use of the reference number without the associated parenthetical is generic to the element. The systems depicted in the accompanying figures are not to scale and components within the figures may be depicted not to scale with each other.

(2) FIG. 1 illustrates a component diagram with an example environment in which application identification may be employed among devices, in accordance with the present concepts.

(3) FIGS. 2 and 3 illustrate example diagrams of communications that devices of FIG. 1 may exchange in order to perform application identification techniques, in accordance with the present concepts.

(4) FIGS. 4 and 5 illustrate flow diagrams of example methods for application identification among devices, in accordance with the present concepts.

(5) FIG. 6 illustrates a computing system diagram illustrating a configuration for a data center that can be utilized to implement aspects of the technologies disclosed herein.

(6) FIG. 7 is a computer architecture diagram showing an illustrative computer hardware architecture for implementing a device that can be utilized to implement aspects of the various technologies presented herein.

DESCRIPTION OF EXAMPLE EMBODIMENTS

(7) Overview

(8) This disclosure describes, at least in part, a method that may be implemented by a client device communicatively coupled to network resources. The method may include receiving an identity request for an identity of an application on a client device. The identity request may be received from the application at a first port of the client device and received by an identification agent at a second port of the client device. The method may include accessing a port list that includes information regarding entities using the first port and the second port of the client device. Based at least in part on the information from the port list, the method may include determining, by the identification agent, the application that is utilizing the first port to communicate the request. The method may include sending, by the identification agent, a query for the identity of the application that is utilizing the first port. The method may also include receiving, in response to the query, the identity of the application. Finally, the method may include sending the identity of the application

to an authentication service on a remote device.

(9) This disclosure also describes, at least in part, a method that may be implemented by network resources communicatively coupled to a client device. The method may include receiving, at an authentication service and from an application on a remote client device, an access request to access a resource managed by the authentication service. The access request may be received via a first communication channel between the authentication service and the remote client device. In response to the access request, the method may include sending a directive to the application. The directive may direct the application to send an identity request for an identity of the application to an identification agent on the remote client device. The method may include receiving, at the authentication service and from the identification agent, the identity of the application, wherein the identity is received via a second communication channel between the authentication service and the remote client device. Based at least in part on the identity of the application, the method may include making, by the authentication service, an access determination regarding whether to allow the application to access to the resource. Further, the method may include sending, to the application at the remote client device, a response to the access request, the response based at least in part on the access determination.

(10) Additionally, the techniques described herein may be performed by a system and/or device having non-transitory computer-readable media storing computer-executable instructions that, when executed by one or more processors, performs the method described above.

EXAMPLE EMBODIMENTS

(11) This disclosure describes techniques for identifying an application (e.g., accessing application) that is attempting to access a resource and/or service, such as a website. In some examples, access may be managed by an authentication service. When an access request is received at the authentication service from an application on a client device, the authentication service may ask the application to communicate with an identification agent on the client device. The identification agent may perform one or more tests regarding the accessing application and/or the client device. Through the one or more tests, the identification agent may discover the identity of the accessing application. In some cases, the identification agent may send the discovered identity to the authentication service. The authentication service may then allow or deny access by the accessing application to the resource based at least in part on the discovered identity.

(12) In some examples, a user may wish to log in to a resource that is available over a network. As used herein, “resource” may represent a wide variety of services or entities that may require authentication and/or authorization to access. For example, a resource may be a website or cloud service with which the user has a user account. The user account may have associated login and/or password information with which the user may authenticate in order to gain access to the user account. For security or management purposes, the resource may establish a relationship with an authentication service purpose to help authenticate and/or authorize users attempting to access the resource. In some cases, an authentication service may be a multiple factor authentication service (e.g., a dual-authentication service). For example, the user may provide the login and/or password information for the user account, and the authentication service may provide an additional level of authentication and/or authorization before determining that a user should be allowed or denied access to the resource. The additional authentication/authorization may include investigating an identity of the application with which the user is attempting to gain access to the resource.

Additionally, the authentication service may learn other information about a client device with which the user is attempting access, collect further information from the user (e.g., biometrics), etc.

(13) For illustration purposes, consider that a user with a client device intends to login to a resource at a remote device. The user utilizes an application on the client device to initiate the login request. The user may enter login and/or password information to attempt to login to a user account. The “application” may be any of a variety of applications that can be used to access a resource. An application could include various types of web browsers, embedded browsers, and/or native

applications on the user device. For example, an application may include a browser running “under the hood,” of the application. As such, the identity or nature of the particular software that is attempting to access the resource may be relatively obscure and/or complicated. The application may not be required to report an identity of the application, or may report some other aspect of the application while the identity of the particular portion that is requesting access remains obscure. Previously, when web browser was being used as an authentication application, the only way to enforce policy around that application was to use the user-agent HTTP header that the embedded browser provided. The techniques described herein makes use of a separate process to map the access request to an identifier that divulges a true identity of the application, outside of what is reported by the user-agent. The example applications described herein are not meant to be limiting.

(14) When the user attempts to login to the resource, the request may be sent to the authentication service at some point in the login process. The request may or may not be routed to or through a device associated with the resource before arriving at the authentication service. Regardless, once the request (or notice of the request) arrives at the authentication service, a communication channel is established between the client device of the user and the authentication service. The authentication service may be located at a particular device (e.g., server) or may potentially be represented by multiple devices. For instance, the authentication service may be cloud-based, providing authentication (and/or authorization, etc.) services through multiple networked cloud computing resources.

(15) In order to make an informed decision regarding allowing or denying access by the client device to the resource, the authentication service may send instructions to the application on the client device to communicate with an identification agent. The identification agent may be software and/or code located on the client device. The identification agent may be a partner and/or extension of the authentication service. In response to the request, the application may communicate with the identification agent. The identification agent may perform one or more tests to assist the authentication service with identification of the accessing application and/or collecting or discovering other information that may help the authentication service to determine whether to allow or deny access to the resource. The tests may include a variety of queries or other methods of acquiring information about the accessing application and/or the client device in general. For example, once communication between the application and the identification agent is established, the identification agent may learn a port (e.g., port ID, port number) through which the application is communicating. The identification agent may use the port ID to look up an identity of the application that is running on that port ID. In this manner, the identification agent may learn the identity of the application without relying on the application to truthfully divulge the identity to either the identification agent or the authentication service. Therefore, the identification agent may be more resistant to spoofing by an application than other methods that rely on information presented or offered by the application (e.g., presented by a user agent).

(16) Once the identification agent has discovered the identity of the application, the identification agent may report the identity and/or other information about the application (e.g., version) and/or about the client device (e.g., operating system) to the authentication service. The identification agent may open a separate channel to communicate with the authentication service. The authentication service may then be able to make a better-informed decision about whether to allow or deny access to the resource. For instance, the authentication service may apply a policy in the decision-making process. The policy may stipulate that certain applications may be used to access the resource, while others should be prevented from being used to access the resource. Since the identity of the application was learned by the identification agent without relying on the application to truthfully divulge the identity, and the identification agent communicated the identity to the authentication service separately (not through the application), the authentication service may be able to more successfully apply the policy. Therefore, security for the resource is improved, through an identification process that is more robust to tricking or fooling, or simply being

presented with incorrect information.

(17) To summarize, a more robust technique for identifying an application that requests access to a resource may improve network security and/or operations. The techniques described herein include using an identification agent located on a client device that is able to discover the identity of an application, without relying on the application to truthfully divulge its identity. The identification agent may be enacted with relatively low overhead and/or low draw on computing resources. In some examples, the identification agent may be viewed as a lightweight mechanism for improving security, featuring both relatively low computational cost and relatively low bandwidth usage. Furthermore, improved decisions regarding allowing or denying access to a resource may help reduce the future cost(s) of inappropriately allowing or denying access, such as the cost to retroactively deny access or resolve access for entities that should have been allowed into the resource.

(18) Although the examples described herein may refer to an identification agent on a client device as the central point of application identification, the techniques can generally be applied to any device in a network. Further, the techniques are generally applicable for any network of devices managed by any entity where virtual resources are provisioned. In some instances, the techniques may be performed by software-defined networking (SDN), and in other examples, various devices may be used in a system to perform the techniques described herein. The devices by which the techniques are performed herein are a matter of implementation, and the techniques described are not limited to any specific architecture or implementation.

(19) The techniques described herein provide various improvements and efficiencies with respect to network communications. For instance, the techniques described herein may reduce the amount of computational resource use, storage, dropped data, latency, and other issues experienced in networks due to lack of network resources, overuse of network resources, issues with timing of network communications, and/or improper routing of data. By improving network communications across a network, overall performance by servers and virtual resources may be improved.

(20) Certain implementations and embodiments of the disclosure will now be described more fully below with reference to the accompanying figures, in which various aspects are shown. However, the various aspects may be implemented in many different forms and should not be construed as limited to the implementations set forth herein. The disclosure encompasses variations of the embodiments, as described herein. Like numbers refer to like elements throughout.

(21) FIG. 1 illustrates an example environment **100** in accordance with the present application identification concepts. Example environment **100** may include network resources **102** (e.g., cloud-based resources), a server device **104**, and/or client devices **106** (e.g., user devices). Network resources **102** may include any of a variety of computing resources such as server devices, applications, databases, storage devices, networks, other computing devices, etc. These computing resources may be viewed as being remote from a client device **106** and/or a server device **104**, and therefore may be referred to simply as one or more remote devices. In FIG. 1, three client devices **106** are depicted, including client device **106(1)**, client device **106(2)**, and client device **106(3)**. Herein, parentheticals are utilized after a reference number to distinguish like elements. Use of the reference number without the associated parenthetical is generic to the element. As such, a client device **106** may represent any of a variety of computing devices of a user, such as a computer, laptop, mobile device, tablet, smart watch, etc.

(22) In example environment **100**, network resources **102**, server device **104**, client devices **106**, and/or other devices may be communicatively coupled to each other via network **108** (e.g., cloud computing network). Within the example environment **100**, server device **104**, network resources **102**, client devices **106**, and/or other devices may exchange communications (e.g., packets) via a network connection(s) to network **108**, indicated by double arrows **110**. For instance, network connections **110** may be transport control protocol (TCP) network connections or any network connection (e.g., information-centric networking (ICN)) that enables the server device **104** to

exchange packets with other devices via network **108**. The network connections **110** represent, for example, data paths between the server device **102**, network resources **104**, and client devices **106**. It should be appreciated that the term “network connection” may also be referred to as a “network path.” Also, the data may pass through one or more other network devices (e.g., router, switch) on any path between the server device **102**, network resources **104**, and/or client devices **106**. The suggestion of a cloud computing network(s) relative to the example environment in FIG. **1** is not meant to be limiting. Other types of networks are contemplated in accordance with application identification concepts.

(23) In some examples, any of client devices **106** may try to access a resource **112** (e.g., website) at server device **104**. For example, a user of client device **106(1)** may wish to log on to a website at server device **104**. Server device **104** may host and/or control access to the website and/or other resource. The client device **106(1)** may use an application **114(1)** to attempt access to the resource **112** at server device **104**. In some cases, the resource **112** may use an authentication service **116** located at network resources **102** to help determine whether to allow or deny access by the application **114(1)**. Further, the authentication service **116** may utilize an identification (ID) agent **118(1)** on client device **106(1)** to help discover information about the application **114(1)** and/or the client device **106(1)**. The authentication service **116** may then be able to use any discovered information, such as an identity of application **114(1)**, to make a decision regarding allowing or denying access to the resource **112**. For instance, the authentication service **116** may apply a policy (e.g., software-based policy) to the identity of application **114(1)** to grant or deny access.

(24) FIG. **1** illustrates example communications between network resources **102**, authentication service **116**, client device **106(1)**, application **114(1)**, and/or identification agent **118(1)** in accordance with the present application identification concepts. The communications are indicated with dashed, numbered lines. For example, the communication(s) at “Step **1**” may include client device **106(1)** using application **114(1)** to communicate with network resources **102** regarding accessing resource **112**. Step **1** may also represent authentication service **116** (and/or network resources **102**) communicating with application **114(1)** to establish communication with identification agent **118(1)**. As a result, at “Step **2**,” application **114(1)** may communicate with identification agent **118(1)**. Also at Step **2**, identification agent **118(1)** may collect information and/or perform tests to help discover information about application **114(1)**. Identification agent may or may not communicate with other entities or components of client device **106(1)** to collect information regarding application **114(1)** or client device **106(1)**. In some examples, when identification agent **118(1)** discovers an identity of application **114(1)** and/or other information about application **114(1)** and/or client device **106(1)**, at “Step **3**” identification agent **118(1)** may communicate any of the information to authentication service **116**. Additional detail regarding potential communications between authentication service **116**, an identification agent **118**, and/or an application **114** will be provided relative to the examples illustrated in FIGS. **2** and **3**, below.

(25) FIG. **2** illustrates an example diagram **200** of communications that entities described relative to FIG. **1** may exchange to perform application identification techniques. For example, diagram **200** may represent a call flow between at least some of the devices and/or applications of FIG. **1**. FIG. **2** includes authentication service **116** (at network resources **102**), an identification agent **118** (at client device **106**), and an application **114** (also at client device **106**).

(26) At **202** of FIG. **2**, authentication service **116** may receive an access request from application **114**. As suggested above, the access request may come more or less directly from client device **106** to authentication service **116**, may be forwarded from another entity (such as the server device **104** that offers the resource **112** (FIG. **1**)), or authentication service **116** may receive some other indication that application **114** on client device **106** wishes to access the resource **112** at server device **104**.

(27) At **204** of FIG. **2**, authentication service **116** may communicate with application **114**. The communication may represent an open channel or other connection between authentication service

116 and application **114**. The communication may include a request and/or other directive for application **114** to communicate with identification agent **118**. For example, authentication service **116** may offer a hypertext markup language (HTML) document to application **114**. The HTML document may include code that directs application **114** to communicate with identification agent **118**, which is also located on client device **106**. For instance, the code may direct application **114** to send a hypertext transfer protocol (HTTP) request, an Asynchronous JavaScript And XML (Ajax) request, or some other communication (e.g., WebSocket) to identification agent **118**.

(28) At **206** of FIG. 2, application **114** may communicate with identification agent **118**. In this example, application **114** sends an Ajax request to identification agent **118**. Application **114** may communicate with identification agent **118** via a first open port within client device **106**. Identification agent **118** may receive the communication via a second open port of client device **106**.

(29) At **208** of FIG. 2, identification agent **118** may continue to communicate with application **114** (and/or with other elements of client device **106**) as part of a collection phase of the application identification process, to collect and/or check information about application **114** and/or other elements of client device **106**. The collection phase may represent one or more tests, queries, and/or responses. Identification agent **118** may use application program interfaces (APIs) to make one or more tests or queries, for instance. In some examples, identification agent **118** may be able to access a list (e.g., port list, table, database) of open ports (e.g., ports in use, active ports) of client device **106**, or otherwise find out which ports of client device **106** are being used for communication, and/or which ports are communicating with each other. For instance, identification agent **118** may query an IPv4 TCP connection table using a loopback IP address as both the remote and local address. Identification agent **118** may further narrow down the results by considering the source and destination ports (first and second ports). In a case where the narrowed results match the result from the IPv4 TCP connection table, the application **114** using the source port (first port) has been located. Identification agent **118** may then make a system call to retrieve a process identification (PID) of the application **114**. Identification agent **118** may use the PID to send a query requesting the name of application **114** using the PID, leaving identification agent **118** with the true identity of application **114** that made the access request. In other examples of the collection phase of the application identification process, identification agent **118** may further learn the version of application **114**, an operating system of client device **106**, whether client device **106** employs a firewall, etc. The information may be related to network and/or device security, general health, operating performance, and/or various other aspects of application **114** and/or client device **106**.

(30) In some cases, aspects and/or an order of a call flow may be based at least in part on an operating system of client device **106**. For purposes of illustration, the description of the collection phase of step **208** provided above may correspond to an example where the operating system of client device **106** is Windows (Microsoft). For comparison, the description of step **208** provided below may correspond to an example where the operating system of client device **106** is macOS (Apple Inc.) or Linux.

(31) At **208** of FIG. 2, in another example, identification agent **118** may retrieve the PID by an alternative method. In this example, identification agent **118** may also make use of system calls to retrieve the PID. However, identification agent **118** may get a list of all running processes by PID. Identification agent **118** may filter out those processes without a PID at all (i.e., the PID is 0 in the query result). Identification agent **118** may further filter down the list by finding a process that has an open communication with a port of the identification agent **118** showing as both the source and destination port, so that the logic arrives at the PID of application **114**. In a case where the application **114** in question is still ambiguous, identification agent **118** may use additional operating system APIs in order to inspect a parent process of application **114**. Additional operating system APIs may allow identification of ambiguous cases. For instance, Apple Safari may initially look

identical to other applications using embedded WebKit views, until an additional inspection is performed.

(32) At **210** of FIG. 2, identification agent **118** may communicate with authentication service **116**. For example, identification agent **118** may send an HTTP request and/or use a HTTPS connection to talk to authentication service **116**. The communication between identification agent **118** and authentication service **116** may represent a second open channel or other connection. In other words, the communication between identification agent **118** and authentication service **116** may represent a second communication channel between client device **106** and network resources **102** that is separate from the first communication channel between authentication service **116** and application **114**. As such, authentication service **116** can trust the information it receives from identification agent **118**, since the information is not routed through a potentially untrusted entity (e.g., application **114**). Authentication service **116** knows to associate the communication from the identification agent **118** with the access request received from application **114**. Identification agent **118** may send information to the authentication service **116**, such as the identity of application **114** and/or other information about application **114** and/or client device **106**. Authentication service **116** is therefore enabled to make a better-informed decision regarding allowing or denying the access request from application **114**. Authentication service **116** may also save the identity or other information. In some examples, authentication service **116** may overwrite previously received identity information, such as an identity reported by a user-agent, since the identity received from identification agent **118** is expected to be more accurate.

(33) FIG. 3 illustrates an additional example diagram **300** of communications that devices of FIG. 1 may exchange to perform application identification techniques. Some aspects of the example communications shown in FIG. 3 may be similar to aspects of the example communications described above relative to FIG. 2. Therefore, for sake of brevity, not all elements of FIG. 3 will be described in detail. For example, steps **302-308** of FIG. 3 may be considered similar to Steps **202-208** of FIG. 2.

(34) FIG. 3 illustrates an alternative method by which the identity of application **114** and/or other information is returned to authentication service **116** from identification agent **118**. In FIG. 3, rather than the identity result being sent directly from identification agent **118** to authentication service **116** by a new channel or other connection, the identity result is passed through application **114**. However, to prevent application **114** from changing the result, for example to hide the identity, identification agent **118** may sign the message so that authentication service **116** knows to trust the identity. At **310** identification agent **118** sends the signed identity result to application **114**. At **312** application **114** forwards the signed identity result to authentication service **116**. Since the message is signed by identification agent **118**, authentication service **116** knows to trust the identity of application **114** contained in the message. Forwarding the identity result through application **114** may prevent the need to establish a separate connection between identification agent **118** and authentication service **116**.

(35) Other variations of application identification techniques are contemplated. For example, some of the decision-making regarding allowing or denying access to the resource could be enabled in identification agent **118**. Stated another way, identification agent **118** could be “smarter,” able to contribute to the decision-making process. A potential benefit could be determining that further information may need to be collected for a well-informed decision to be made, and identification agent **118** could initiate collecting additional, targeted information. In another example variation, identification agent **118** could monitor various applications and/or software used at client device **106** over time. When an access request is presented, identification agent **118** could determine whether the requesting application is regularly used at client device **106**, or unusual for use at client device **106**. A policy governing allowing or denying access may be written to specify that an unusual application may be considered more suspicious, and may indicate that the access should be denied.

(36) To summarize, the application identification techniques described herein may enable accurate identification of software that is attempting access a resource. The techniques may be relatively lightweight, featuring low computational cost and/or low bandwidth usage. Even when an HTTP user-agent reports indistinguishable and/or erroneous values, the present techniques utilize operating system APIs to accurately identify the requesting software. After the identity collection phase, which is swift and not system-resource intensive, the information may be relayed back to authentication service which can then more accurately apply access policy before allowing a user to authenticate. More appropriate control of access to resources can help prevent further undue computational cost and/or bandwidth usage, improving performance of the entities involved. Therefore, application identification techniques may improve security and/or performance of network resources.

(37) FIGS. 4 and 5 illustrate flow diagrams of example methods **400** and **500** that include functions that may be performed at least partly by a client device and/or network resources, such as client devices **106** and/or network resources **102** described relative to FIG. 1. The logical operations described herein with respect to FIGS. 4 and 5 may be implemented (1) as a sequence of computer-implemented acts or program modules running on a computing system and/or (2) as interconnected machine logic circuits or circuit modules within the computing system.

(38) The implementation of the various devices and/or components described herein is a matter of choice dependent on the performance and other requirements of the computing system. Accordingly, the logical operations described herein are referred to variously as operations, structural devices, acts, or modules. These operations, structural devices, acts, and modules may be implemented in software, in firmware, in special purpose digital logic, and any combination thereof. It should also be appreciated that more or fewer operations might be performed than shown in the FIGS. 4 and 5 and described herein. These operations may also be performed in parallel, or in a different order than those described herein. Some or all of these operations may also be performed by components other than those specifically identified. Although the techniques described in this disclosure are with reference to specific devices, in other examples, the techniques may be implemented by less devices, more devices, different devices, or any configuration of devices and/or components.

(39) FIG. 4 illustrates a flow diagram of an example method **400** for a network device to perform application identification techniques. Method **400** may be performed by a client device (e.g., client device **106**) communicatively coupled to network resources (e.g., network resources **102**) and/or other devices, for instance. In some examples, method **400** may be performed by a computing device comprising one or more processors and one or more non-transitory computer-readable media storing computer-executable instructions that, when executed by the one or more processors, cause the one or more processors to perform method **400**.

(40) At **402**, method **400** may include receiving an identity request for an identity of an application on a client device. In some examples, the identity request may be received from the application at a first port of the client device. Also, the identity request may be received by an identification agent at a second port of the client device. The identity request may be received at the application in response to an earlier action by the application. For example, the application may have initiated communication with an authentication service at a remote device by sending an access request to the authentication service. The access request may correspond to a resource managed by the authentication service. For instance, the authentication service may manage access to the resource for a separate server device that hosts or owns the resource. In some examples, in response to the access request, the authentication service may send a directive to the application, directing the application to send the identity request for the identity of the application to the identification agent on the client device. The authentication service may be interested in the identity of the application in order to determine whether to grant or deny access to the resource, for instance. The identity request for the identity of the application may be an HTTP request, an Ajax request, and/or some

other form of communication request and/or invitation sent from the application to the identification agent, in compliance with the directive from the authentication service, for instance. (41) At **404**, method **400** may include accessing a port list that includes information regarding entities that may be using the first port and the second port of the client device. In some examples, the port list may be a transport control protocol (TCP) connection table that includes information regarding which ports of the client device are in use, are connected in communication, and/or which entity is using any given port.

(42) At **406**, based at least in part on the information from the port list, method **400** may include determining, by the identification agent, the application that is utilizing the first port to communicate the request. In some examples, the identification agent may initiate a system call. In response to the system call, the identification agent may receive a process identification (PID) of the application. The PID may be used by the identification agent to gather further information regarding the application.

(43) At **408**, method **400** may include sending, by the identification agent, a query for the identity of the application that is utilizing the first port. In some examples, the query may be based on the PID of the application received in response to the system call.

(44) At **410**, in response to the query, method **400** may include receiving the identity of the application. The identity of the application may include a browser identity of a browser. The browser may have been used to communicate between the application and the authentication service. For instance, the browser may have sent the access request to the authentication service.

(45) At **412**, method **400** may include sending the identity of the application to the authentication service, which may be located on a remote device. The identity of the application sent to the authentication service may include the browser identity, a version of the browser, and/or other information regarding the application or the client device, such as an operating system of the client device, etc.

(46) FIG. 5 illustrates a flow diagram of an example method **500** for network devices to perform application identification techniques. Method **400** may be performed by network resources (e.g., network resources **102**) communicatively coupled to a client device (e.g., client device **106**) and/or other devices, for instance. In some examples, method **500** may be performed by a computing device comprising one or more processors and one or more non-transitory computer-readable media storing computer-executable instructions that, when executed by the one or more processors, cause the one or more processors to perform method **500**.

(47) At **502**, method **500** may include receiving, at an authentication service offered via the network resources and from an application on a remote client device, an access request. The access request may indicate that a user of the client device wishes to access a resource managed by the authentication service. In some examples, the access request may be received via a first communication channel between the authentication service and the remote client device.

(48) At **504**, in response to the access request, method **500** may include sending a directive to the application. The directive may direct the application to send an identity request for an identity of the application to an identification agent on the remote client device.

(49) At **506**, method **500** may include receiving, at the authentication service and from the identification agent, the identity of the application. The identity of the application may include a browser identity of a browser that is used to communicate between the remote client device and the authentication service. The identity may include additional information in some cases, such as a version of the browser. In some examples, the identity may be received via a second communication channel between the authentication service and the remote client device.

(50) At **508**, based at least in part on the identity of the application, method **500** may include having the authentication service make an access determination regarding whether to allow the application to access to the resource. In some examples, the access determination may be performed by applying a policy. For instance, the policy may consider the identity of the

application in determining whether to allow or deny access to the resource by the application. Additionally, the authentication service may store the identity of the application within the network resources, which may be useful to the authentication service for future reference. In some cases, the authentication service may overwrite previously received identity information corresponding to the application. For instance, a previously received identity of the application may have been received from a user-agent associated with the application, and therefore may be less trusted than the identity received from the identification agent.

(51) At **510**, method **500** may include sending a response to the access request from the authentication service to the application at the remote client device based at least in part on the access determination. For instance, the response may include an indication that access to the resource is allowed or denied. In some examples, the response may include further communication regarding a login process for the client device to access the resource. In other examples, the authentication service may send the indication that access to the resource is allowed or denied to a server device that hosts the resource, rather than to the client device. Additionally or alternatively, the authentication service may simply proceed with facilitating access by the client device to the resource, without sending an indication of the access determination.

(52) FIG. **6** is a computing system diagram illustrating a configuration for a data center **600** that can be utilized to implement aspects of the technologies disclosed herein. The example data center **600** shown in FIG. **6** includes several computers **602A-602F** (which might be referred to herein singularly as “a computer **602**” or in the plural as “the computers **602**”) for providing computing resources. In some examples, the resources and/or computers **602** may include, or correspond to, any type of networked device described herein, such as a network resources (**102**), network **108**, and/or server device (**104**). Although, computers **602** may comprise any type of networked device, such as servers, switches, routers, hubs, bridges, gateways, modems, repeaters, access points, hosts, etc.

(53) The computers **602** can be standard tower, rack-mount, or blade server computers configured appropriately for providing computing resources. In some examples, the computers **602** may provide computing resources **604** including data processing resources such as virtual machine (VM) instances or hardware computing systems, database clusters, computing clusters, storage clusters, data storage resources, database resources, networking resources, and others. Some of the computers **602** can also be configured to execute a resource manager **606** capable of instantiating and/or managing the computing resources. In the case of VM instances, for example, the resource manager **606** can be a hypervisor or another type of program configured to enable the execution of multiple VM instances on a single computer **602**. Computers **602** in the data center **600** can also be configured to provide network services and other types of services.

(54) In the example data center **600** shown in FIG. **6**, an appropriate local area network (LAN) **608** is also utilized to interconnect the computers **602A-602F**. It should be appreciated that the configuration and network topology described herein has been greatly simplified and that many more computing systems, software components, networks, and networking devices can be utilized to interconnect the various computing systems disclosed herein and to provide the functionality described above. Appropriate load balancing devices or other types of network infrastructure components can also be utilized for balancing a load between data centers **600**, between each of the computers **602A-602F** in each data center **600**, and, potentially, between computing resources in each of the computers **602**. It should be appreciated that the configuration of the data center **600** described with reference to FIG. **6** is merely illustrative and that other implementations can be utilized.

(55) In some examples, the computers **602** may each execute one or more application containers and/or virtual machines to perform techniques described herein. For instance, the containers and/or virtual machines may serve as server devices, user devices, and/or routers in network **108** and/or network resources **102**.

(56) In some instances, the data center **600** may provide computing resources, like application containers, VM instances, and storage, on a permanent or an as-needed basis. Among other types of functionality, the computing resources provided by a cloud computing network may be utilized to implement the various services and techniques described above. The computing resources **604** provided by the cloud computing network can include various types of computing resources, such as data processing resources like application containers and VM instances, data storage resources, networking resources, data communication resources, network services, and the like.

(57) Each type of computing resource **604** provided by the cloud computing network can be general-purpose or can be available in a number of specific configurations. For example, data processing resources can be available as physical computers or VM instances in a number of different configurations. The VM instances can be configured to execute applications, including web servers, application servers, media servers, database servers, some or all of the network services described above, and/or other types of programs. Data storage resources can include file storage devices, block storage devices, and the like. The cloud computing network can also be configured to provide other types of computing resources **604** not mentioned specifically herein.

(58) The computing resources **604** provided by a cloud computing network may be enabled in one embodiment by one or more data centers **600** (which might be referred to herein singularly as “a data center **600**” or in the plural as “the data centers **600**”). The data centers **600** are facilities utilized to house and operate computer systems and associated components. The data centers **600** typically include redundant and backup power, communications, cooling, and security systems. The data centers **600** can also be located in geographically disparate locations. One illustrative embodiment for a data center **600** that can be utilized to implement the technologies disclosed herein will be described below with regard to FIG. 7.

(59) FIG. 7 shows an example computer architecture **700** for a computer **602** capable of executing program components for implementing the functionality described above. The computer architecture **700** shown in FIG. 7 illustrates a conventional server computer, workstation, desktop computer, laptop, tablet, network appliance, e-reader, smartphone, and/or other computing device, and can be utilized to execute any of the software components presented herein. The computer **602** may, in some examples, correspond to a physical device described herein (e.g., server device, user device, router, etc.), and may comprise networked devices such as servers, switches, routers, hubs, bridges, gateways, modems, repeaters, access points, etc. For instance, computer **602** may correspond to client device **106**, network resources **102**, and/or server device **104**.

(60) As shown in FIG. 7, the computer **602** includes a baseboard **702**, or “motherboard,” which is a printed circuit board to which a multitude of components or devices can be connected by way of a system bus or other electrical communication paths. In one illustrative configuration, one or more central processing units (“CPUs”) **704** operate in conjunction with a chipset **706**. The CPUs **704** can be standard programmable processors that perform arithmetic and logical operations necessary for the operation of the computer **602**.

(61) The CPUs **704** perform operations by transitioning from one discrete, physical state to the next through the manipulation of switching elements that differentiate between and change these states. Switching elements generally include electronic circuits that maintain one of two binary states, such as flip-flops, and electronic circuits that provide an output state based on the logical combination of the states of one or more other switching elements, such as logic gates. These basic switching elements can be combined to create more complex logic circuits, including registers, adders-subtractors, arithmetic logic units, floating-point units, and the like.

(62) The chipset **706** provides an interface between the CPUs **704** and the remainder of the components and devices on the baseboard **702**. The chipset **706** can provide an interface to a RAM **708**, used as the main memory in the computer **602**. The chipset **706** can further provide an interface to a computer-readable storage medium such as a read-only memory (“ROM”) **710** or non-volatile RAM (“NVRAM”) for storing basic routines that help to startup the computer **602** and

to transfer information between the various components and devices. The ROM **710** or NVRAM can also store other software components necessary for the operation of the computer **602** in accordance with the configurations described herein.

(63) The computer **602** can operate in a networked environment using logical connections to remote computing devices and computer systems through a network, such as the network **108** and/or **608**. The chipset **706** can include functionality for providing network connectivity through a network interface controller (NIC) **712**, such as a gigabit Ethernet adapter. The NIC **712** is capable of connecting the computer **602** to other computing devices over the network **108**. For instance, in the example shown in FIG. 7, NIC **712** may help facilitate transfer of data, packets, and/or communications over the network **108** between client device **106** and network resources **102**. It should be appreciated that multiple NICs **712** can be present in the computer **602**, connecting the computer to other types of networks and remote computer systems.

(64) The computer **602** can be connected to a storage device **714** that provides non-volatile storage for the computer. The storage device **714** can store an operating system **716**, programs **718**, and/or other data, such as a port list, TCP connection table, application identity information, etc. The storage device **714** can be connected to the computer **602** through a storage controller **722** connected to the chipset **706**, for example. The storage device **714** can consist of one or more physical storage units. The storage controller **722** can interface with the physical storage units through a serial attached SCSI (“SAS”) interface, a serial advanced technology attachment (“SATA”) interface, a fiber channel (“FC”) interface, or other type of interface for physically connecting and transferring data between computers and physical storage units.

(65) The computer **602** can store data on the storage device **714** by transforming the physical state of the physical storage units to reflect the information being stored. The specific transformation of physical state can depend on various factors, in different embodiments of this description. Examples of such factors can include, but are not limited to, the technology used to implement the physical storage units, whether the storage device **714** is characterized as primary or secondary storage, and the like.

(66) For example, the computer **602** can store information to the storage device **714** by issuing instructions through the storage controller **722** to alter the magnetic characteristics of a particular location within a magnetic disk drive unit, the reflective or refractive characteristics of a particular location in an optical storage unit, or the electrical characteristics of a particular capacitor, transistor, or other discrete component in a solid-state storage unit. Other transformations of physical media are possible without departing from the scope and spirit of the present description, with the foregoing examples provided only to facilitate this description. The computer **602** can further read information from the storage device **714** by detecting the physical states or characteristics of one or more particular locations within the physical storage units.

(67) In addition to the mass storage device **714** described above, the computer **602** can have access to other computer-readable storage media to store and retrieve information, such as policies, program modules, data structures, and/or other data. It should be appreciated by those skilled in the art that computer-readable storage media is any available media that provides for the non-transitory storage of data and that can be accessed by the computer **602**. In some examples, the operations performed by the network **108**, and or any components included therein, may be supported by one or more devices similar to computer **602**. Stated otherwise, some or all of the operations performed by the network **108**, and or any components included therein, may be performed by one or more computer devices **602** operating in a cloud-based arrangement.

(68) By way of example, and not limitation, computer-readable storage media can include volatile and non-volatile, removable and non-removable media implemented in any method or technology. Computer-readable storage media includes, but is not limited to, RAM, ROM, erasable programmable ROM (“EPROM”), electrically-erasable programmable ROM (“EEPROM”), flash memory or other solid-state memory technology, compact disc ROM (“CD-ROM”), digital

versatile disk (“DVD”), high definition DVD (“HD-DVD”), BLU-RAY, ternary content addressable memory (TCAM), and/or other optical storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium that can be used to store the desired information in a non-transitory fashion.

(69) As mentioned briefly above, the storage device **714** can store an operating system **716** utilized to control the operation of the computer **602**. According to one embodiment, the operating system comprises the LINUX operating system. According to another embodiment, the operating system comprises the WINDOWS® SERVER operating system from MICROSOFT Corporation of Redmond, Washington. According to further embodiments, the operating system can comprise the UNIX operating system or one of its variants. It should be appreciated that other operating systems can also be utilized. The storage device **714** can store other system or application programs and data utilized by the computer **602**.

(70) In one embodiment, the storage device **714** or other computer-readable storage media is encoded with computer-executable instructions which, when loaded into the computer **602**, transform the computer from a general-purpose computing system into a special-purpose computer capable of implementing the embodiments described herein. These computer-executable instructions transform the computer **602** by specifying how the CPUs **704** transition between states, as described above. According to one embodiment, the computer **602** has access to computer-readable storage media storing computer-executable instructions which, when executed by the computer **602**, perform the various processes described above with regard to FIGS. 1-5. The computer **602** can also include computer-readable storage media having instructions stored thereupon for performing any of the other computer-implemented operations described herein.

(71) The computer **602** can also include one or more input/output controllers **724** for receiving and processing input from a number of input devices, such as a keyboard, a mouse, a touchpad, a touch screen, an electronic stylus, or other type of input device. Similarly, an input/output controller **724** can provide output to a display, such as a computer monitor, a flat-panel display, a digital projector, a printer, or other type of output device. It will be appreciated that the computer **602** might not include all of the components shown in FIG. 7, can include other components that are not explicitly shown in FIG. 7, or might utilize an architecture completely different than that shown in FIG. 7.

(72) As described herein, the computer **602** may comprise one or more devices, such as client device **106**, server device **104**, network resources **106**, and/or other devices. The computer **602** may include one or more hardware processors **704** (processors) configured to execute one or more stored instructions. The processor(s) **704** may comprise one or more cores. Further, the computer **602** may include one or more network interfaces configured to provide communications between the computer **602** and other devices, such as the communications described herein as being performed by client device **106**, network resources **102**, and/or other devices. In some examples, the communications may include data, packet(s), request(s), response(s), and/or other information transfer, for instance. The network interfaces may include devices configured to couple to personal area networks (PANs), wired and wireless local area networks (LANs), wired and wireless wide area networks (WANs), and so forth. For example, the network interfaces may include devices compatible with Ethernet, Wi-Fi™, and so forth.

(73) The programs **718** may comprise any type of programs or processes to perform the techniques described in this disclosure in accordance with application identification techniques. For instance, the programs **718** may cause the computer **602** to perform techniques for communicating with other devices or among elements of a device using any type of protocol or standard usable for determining connectivity. Additionally, the programs **718** may comprise instructions that cause the computer **602** to perform the specific techniques for application identification.

(74) While the invention is described with respect to the specific examples, it is to be understood that the scope of the invention is not limited to these specific examples. Since other modifications and changes varied to fit particular operating requirements and environments will be apparent to

those skilled in the art, the invention is not considered limited to the example chosen for purposes of disclosure, and covers all changes and modifications which do not constitute departures from the true spirit and scope of this invention.

(75) Although the application describes embodiments having specific structural features and/or methodological acts, it is to be understood that the claims are not necessarily limited to the specific features or acts described. Rather, the specific features and acts are merely illustrative some embodiments that fall within the scope of the claims of the application.

Claims

1. A method performed by an authentication service, the method comprising: receiving, from an application on a remote client device, an access request to access a resource managed by the authentication service, the access request received via a first communication channel between the authentication service and the remote client device; in response to the access request, sending a directive to the application, the directive directing the application to send an identity request for an identity of the application to an identification agent on the remote client device, the identification agent configured to, in response to the directive, make an identity determination of an identity of the application without relying on the application for the identity determination; receiving, from the identification agent, the identity of the application, wherein the identity is received via a second communication channel between the authentication service and the identification agent, the second communication channel being separate from the first communication channel; based at least in part on the identity of the application, making an access determination regarding whether to allow the application to access to the resource; and sending, to the application at the remote client device, a response to the access request, the response based at least in part on the access determination.
2. The method of claim 1, further comprising: applying a policy in order to make the access determination, wherein the identity of the application is used by the policy to make the access determination to allow or deny access to the resource by the application.
3. The method of claim 1, further comprising: overwriting, by the authentication service, a previously received identity of the application, wherein the previously received identity of the application was received from a user-agent associated with the application.
4. The method of claim 1, wherein the identity of the application includes a browser identity and a version of a browser used to communicate between the remote client device and the authentication service.
5. One or more network devices comprising: one or more processors; and one or more non-transitory computer-readable media storing computer-executable instructions that, when executed by the one or more processors, cause the one or more processors to: receive, from an application on a remote client device, an access request to access a resource managed by an authentication service at the one or more network devices, the access request received via a first communication channel between the authentication service and the remote client device; in response to the access request, send a directive to the application, the directive directing the application to send an identity request for an identity of the application to an identification agent on the remote client device, the identification agent configured to, in response to the directive, make an identity determination of an identity of the application without relying on the application for the identity determination; receive, from the identification agent, the identity of the application, wherein the identity is received via a second communication channel between the authentication service and the identification agent, the second communication channel being separate from the first communication channel; based at least in part on the identity of the application, make an access determination regarding whether to allow the application to access to the resource; and send, to the application at the remote client device, a response to the access request, the response based at least in part on the access determination.
6. The one or more network devices of claim 5, wherein the computer-executable instructions

further cause the one or more processors to: apply a policy in order to make the access determination, wherein the identity of the application is used with the policy to make the access determination to allow or deny access to the resource by the application.

7. The one or more network devices of claim 5, wherein the computer-executable instructions further cause the one or more processors to: store the identity of the application received via the second communication channel between the authentication service and the remote client device.

8. The one or more network devices of claim 7, wherein the computer-executable instructions further cause the one or more processors to: overwrite, by the authentication service, a previously received identity of the application, wherein the previously received identity of the application was received from a user-agent associated with the application.

9. The one or more network devices of claim 5, wherein the identity of the application includes a browser identity and a version of a browser used to communicate between the remote client device and the authentication service.

10. The one or more network devices of claim 5, wherein making the identity determination comprises accessing identity information related to the application and further wherein receiving the identity comprises receiving the identity information.

11. The one or more network devices of claim 5, wherein the computer-executable instructions further cause the one or more processors to: receive, by the authentication service and from the identification agent, an indication of whether to allow the application to access the resource.

12. The one or more network devices of claim 11, wherein the indication is used with the identity to make the access determination to allow or deny access to the resource by the application.

13. One or more non-transitory computer-readable media storing instructions that, when executed, cause one or more processors to perform operations of an authentication service, the operations comprising: receiving, from an application on a remote client device, an access request to access a resource managed by the authentication service, the access request received via a first communication channel between the authentication service and the remote client device; in response to the access request, sending a directive to the application, the directive directing the application to send an identity request for an identity of the application to an identification agent on the remote client device, the identification agent configured to, in response to the directive, make an identity determination of an identity of the application without relying on the application for the identity determination; receiving, from the identification agent, the identity of the application, wherein the identity is received via a second communication channel between the authentication service and the identification agent, the second communication channel being separate from the first communication channel; based at least in part on the identity of the application, making an access determination regarding whether to allow the application to access to the resource; and sending, to the application at the remote client device, a response to the access request, the response based at least in part on the access determination.

14. The one or more non-transitory computer-readable media of claim 13, the operations further comprising: applying a policy in order to make the access determination, wherein the identity of the application is used by the policy to make the access determination to allow or deny access to the resource by the application.

15. The one or more non-transitory computer-readable media of claim 13, the operations further comprising: storing the identity of the application received via the second communication channel between the authentication service and the remote client device.

16. The one or more non-transitory computer-readable media of claim 15, the operations further comprising: overwriting a previously received identity of the application, wherein the previously received identity of the application was received from a user-agent associated with the application.

17. The one or more non-transitory computer-readable media of claim 13, wherein the identity of the application includes a browser identity and a version of a browser used to communicate between the remote client device and the authentication service.

18. The one or more non-transitory computer-readable media of claim 13, wherein making the identity determination comprises accessing identity information related to the application and further wherein receiving the identity comprises receiving the identity information.

19. The one or more non-transitory computer-readable media of claim 13, the operations further comprising: receiving, from the identification agent, an indication of whether to allow the application to access the resource.

20. The one or more non-transitory computer-readable media of claim 19, wherein the indication is used with the identity to make the access determination to allow or deny access to the resource by the application.
