

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12413621
Kind Code	B2
Date of Patent	September 09, 2025
Inventor(s)	N; Shashi KIran

Visual detection of phishing websites via headless browser

Abstract

There is disclosed a computer-implemented system and method for protecting a user from phishing attacks, including detecting that the user has opened a website associated with a uniform resource locator (URL), and determining that the URL does not have a known good reputation; rendering a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspecting the test version of the website with a digital eye, and determining that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warning the user that the website is or may be a phishing website.

Inventors:	N; Shashi KIran (Bangalore, IN)
Applicant:	McAfee, LLC (San Jose, CA)
Family ID:	73172576
Assignee:	McAfee, LLC (San Jose, CA)
Appl. No.:	18/649734
Filed:	April 29, 2024

Prior Publication Data

Document Identifier	Publication Date
US 20240283821 A1	Aug. 22, 2024

Related U.S. Application Data

continuation parent-doc US 17882460 20220805 US 12069091 child-doc US 18649734
continuation parent-doc US 16676939 20191107 US 11411992 20220809 child-doc US 17882460

Publication Classification

Int. Cl.: H04L9/40 (20220101); G06F16/955 (20190101); G06N20/00 (20190101)

U.S. Cl.:

CPC H04L63/1483 (20130101); G06F16/955 (20190101); G06N20/00 (20190101);
H04L63/1416 (20130101); H04L63/1433 (20130101); H04L63/166 (20130101);

Field of Classification Search

CPC: H04L (63/1483); H04L (63/166); G06F (16/955); G06N (20/00)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7890612	12/2010	Todd et al.	N/A	N/A
8332947	12/2011	Bregman	726/25	H04L 63/1433
8806622	12/2013	Waterson et al.	N/A	N/A
8819826	12/2013	Sallam	N/A	N/A
9009820	12/2014	McDougal et al.	N/A	N/A
9516055	12/2015	Liu	N/A	N/A
9516058	12/2015	Antonakakis et al.	N/A	N/A
9621566	12/2016	Gupta et al.	N/A	N/A
10044739	12/2017	Muttik	N/A	N/A
10574696	12/2019	Celik	N/A	N/A
10581883	12/2019	Syme et al.	N/A	N/A
10834128	12/2019	Rajagopalan	N/A	H04L 51/212
10999322	12/2020	Yuan	N/A	G06F 16/583
11003775	12/2020	Kraemer et al.	N/A	N/A
11146576	12/2020	Mushtaq	N/A	H04L 63/1416
11381597	12/2021	Lancioni et al.	N/A	N/A
2006/0253578	12/2005	Dixon	709/225	G06F 21/44
2006/0253580	12/2005	Dixon	709/225	H04L 63/0227
2007/0118528	12/2006	Choi et al.	N/A	N/A
2007/0245422	12/2006	Hwang et al.	N/A	N/A
2007/0283000	12/2006	Proux	709/224	H04L 63/1416
2007/0294352	12/2006	Shraim et al.	N/A	N/A
2008/0133540	12/2007	Hubbard et al.	N/A	N/A
2010/0100958	12/2009	Jeremiah	N/A	N/A
2010/0332593	12/2009	Barash et al.	N/A	N/A
2012/0158626	12/2011	Zhu et al.	N/A	N/A
2013/0074185	12/2012	McDougal et al.	N/A	N/A
2013/0227636	12/2012	Bettini et al.	N/A	N/A

2014/0359760	12/2013	Gupta et al.	N/A	N/A
2017/0155665	12/2016	Dufour	N/A	G06F 21/566
2017/0195363	12/2016	Dahan et al.	N/A	N/A
2018/0027013	12/2017	Wright et al.	N/A	N/A
2018/0054737	12/2017	Guo et al.	N/A	N/A
2018/0191778	12/2017	Volkov	N/A	N/A
2019/0014149	12/2018	Cleveland et al.	N/A	N/A
2019/0068638	12/2018	Bartik et al.	N/A	N/A
2019/0327267	12/2018	Onut	N/A	H04L 67/02
2020/0036751	12/2019	Kohavi	N/A	N/A
2020/0252428	12/2019	Gardezi et al.	N/A	N/A
2021/0014269	12/2020	Devane	N/A	H04L 63/1483
2021/0021638	12/2020	Lancioni et al.	N/A	N/A
2021/0037006	12/2020	Belenko	N/A	G06F 16/955
2021/0075826	12/2020	Johnson	N/A	H04L 67/02

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
3599753	12/2019	EP	N/A
1020080027035	12/2007	KR	N/A
2016034935	12/2015	WO	N/A
2021016142	12/2020	WO	N/A

OTHER PUBLICATIONS

Cordero, et al., “Catching Phish: Detecting Phishing Attacks from Rendered Website Images,” Dec. 2006. cited by applicant

Primary Examiner: Song; Hee K

Attorney, Agent or Firm: Patent Capital Group

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a continuation (and claims the benefit of priority under 35 U.S.C. § 120) of U.S. application Ser. No. 17/882,460, filed Aug. 5, 2022, and entitled “Visual Detection of Phishing Websites Via Headless Browser,” which is a continuation of U.S. application Ser. No. 16/676,939, filed Nov. 7, 2019 and entitled “Visual Detection of Phishing Websites Via Headless Browser.” The disclosures of the prior applications are considered part of and are incorporated by reference in the disclosure of this Application.

FIELD OF THE SPECIFICATION

(1) This application relates in general to computer security, and more particularly, though not exclusively, to providing a system and method for visual detection of phishing websites via headless browser.

BACKGROUND

(2) Modern computing ecosystems often include “always on” broadband internet connections. These connections leave computing devices exposed to the internet, and the devices may be vulnerable to attack.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) The present disclosure is best understood from the following detailed description when read with the accompanying FIGURES. It is emphasized that, in accordance with the standard practice in the industry, various features are not necessarily drawn to scale, and are used for illustration purposes only. Where a scale is shown, explicitly or implicitly, it provides only one illustrative example. In other embodiments, the dimensions of the various features may be arbitrarily increased or reduced for clarity of discussion. Furthermore, the various block diagrams illustrated herein disclose only one illustrative arrangement of logical elements. Those elements may be rearranged in different configurations, and elements shown in one block may, in appropriate circumstances, be moved to a different block or configuration.

(2) FIG. 1 is a block diagram of selected elements of a security ecosystem.

(3) FIG. 2 is a block diagram illustrating mitigation of an example phishing attack.

(4) FIG. 3 is a block diagram illustrating a zero day phishing analysis pipeline.

(5) FIG. 4 is a block diagram illustrating the use of a phishing analysis pipeline in an enterprise e-mail server.

(6) FIG. 5 is an illustration of an enterprise gateway.

(7) FIGS. 6-7 are a flowchart of a method of performing zero day phishing analysis.

(8) FIG. 8 is a block diagram illustrating selected elements of a hardware platform.

(9) FIG. 9 is a block diagram illustrating selected elements of a system-on-a-chip (SoC).

(10) FIG. 10 is a block diagram of selected elements of a processor.

(11) FIG. 11 is a block diagram of selected elements of a network function virtualization (NFV) infrastructure.

SUMMARY

(12) There is disclosed a computer-implemented system and method for protecting a user from phishing attacks, including detecting that the user has opened a website associated with a uniform resource locator (URL), and determining that the URL does not have a known good reputation; rendering a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspecting the test version of the website with a digital eye, and determining that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warning the user that the website is or may be a phishing website.

Embodiments of the Disclosure

(13) The following disclosure provides many different embodiments, or examples, for implementing different features of the present disclosure. Specific examples of components and arrangements are described below to simplify the present disclosure. These are, of course, merely examples and are not intended to be limiting. Further, the present disclosure may repeat reference numerals and/or letters in the various examples. This repetition is for the purpose of simplicity and clarity and does not in itself dictate a relationship between the various embodiments and/or configurations discussed. Different embodiments may have different advantages, and no particular advantage is necessarily required of any embodiment.

(14) Phishing is an important threat to both enterprise and home computer users. Phishing is the use of targeted e-mails and/or webpages that try to induce users into providing personal information, such as usernames, passwords, credit card data, Social Security number, tax information, or other sensitive data. In an illustrative phishing attack, the attacker sends out a mass e-mail blast to a large number of e-mail addresses. While traditional e-mail spam is merely annoying, phishing e-mails are more dangerous. The phishing e-mail may use logos, fonts, backgrounds, and other visual elements selected from popular websites and services. This creates a very sophisticated visual

decoy that superficially appears to originate from the actual service. For example, a user may receive an unsolicited e-mail that appears to be from Netflix. The e-mail may “warn” the user that his Netflix account has a problem, and he needs to log back in and correct his credit card information to resume service.

(15) If the user clicks on the link in the e-mail, he is directed to a phishing website. The phishing website may closely mimic the Netflix login page, and when the user enters his Netflix credentials, the attacker has now collected a valid Netflix login. Even worse, the user is directed to a page that closely mimics the Netflix payment information page. This page indicates that the user's credit card has expired, or that there has been an issue, and requests that the user enter a valid credit card. When the user enters a credit card, the attacker now has a valid credit card number, password, and security code that can be used for identity theft. This identity theft can continue until the user discovers the problem and cancels the credit card.

(16) Other attacks are targeted at accessing online banking credentials, Social Security numbers, and other sensitive data that can be used for identity theft, or for theft of financial services.

(17) In the early days of the internet, phishing attacks were less sophisticated, and only the highly credulous were generally susceptible to them. However, as users became more savvy, phishing attacks became more sophisticated, and it can now be very difficult for even a skeptical user to detect phishing e-mails, or to separate phishing e-mails from non-phishing e-mails.

(18) Computer security mechanisms have gotten to be good at identifying phishing e-mails and websites when they are encountered and scanned. But part of the evolution of phishing attacks is that domain names and URLs used for phishing are rapidly moving targets. In practice, the time to live for a phishing website may be on the order of approximately four hours before the URL is changed. In contrast, web crawlers provided by search engines such as Google may take on the order of two days to scan and categorize the large number of new websites (numbering in the millions or billions) that come online each day. Thus, in many cases, by the time a phishing website is identified, scanned, and categorized, it is already stale and has moved to a new URL.

(19) Human users are, therefore, often tasked with being the first line of defense against phishing scams. However, human users are generally poor at identifying phishing websites. The visual similarity to a trusted website may create a false sense of security in the users. Furthermore, phishing websites often use deceptively similar URLs. For example, the correct, full, unqualified domain name of a legitimate website may be used as a subdomain to a phishing domain. Thus, when the user looks at the link address, he may see what appears to be the correct domain name of a trusted website, followed by a large string of alphanumeric characters that are visually difficult to distinguish from HTML GET code.

(20) Furthermore, the phishing e-mails may be contextually credible. It is not inconceivable to the user that his credit card may have expired, or that there may have been some other problem with his Netflix account, bank account, or some other online service. Indeed, online services may in reality provide such e-mails as a courtesy service to their users when an actual issue of this type arises. Well-crafted phishing e-mails may, in fact, be based on and closely mimic such legitimate notifications, with the exception that they redirect the user to a phishing scam website, instead of to the online service's legitimate website.

(21) Thus, human users often remain the weakest link in the current and evolving threat landscape. In particular, e-mail users are often the soft target of choice for cyber criminals.

(22) Furthermore, phishing attacks remain one of the easiest and most preferable ways for cyber criminals to acquire sensitive data. While modern security infrastructure makes traditional “brute force” electronic attacks difficult for all but the most sophisticated attackers, social engineering attacks like phishing scams continue to be relatively simple and effective. Indeed, contemporary security data indicate that some 80 to 90 percent of cyber attacks include phishing attacks in at least part of the attack lifecycle.

(23) The threat may be even more substantial for enterprises. One study shows that approximately

97% of people worldwide are unable to successfully identify a sophisticated phishing e-mail or website.

(24) Thus, even with sophisticated technology infrastructure, enterprises and consumers are still vulnerable to well-targeted phishing attacks, unless and until a savvy user identifies the malicious URL and has it blocked.

(25) While the visual similarity of phishing e-mails and websites is a key enabler for phishing attacks from the perspective of human users, that very visual similarity can be an Achilles' heel for computer analysis. And while human users have relative difficulty in identifying phishing websites and e-mails, computer analysis is relatively straightforward.

(26) Therefore, one consideration is to subject a suspicious website to sophisticated computer analysis before human users even have a chance to act on it. As a practical matter, it may not be preferable to scan every single URL contained within an e-mail or visited by a user. To do so may make the browsing experience unacceptably slow to users, and inhibit legitimate uses and business operations. The use of intelligent filtering can help to minimize the number of URLs submitted to scanning. For example, in practice, on the order of 95 percent of URLs most users visit may have a known reputation, either a globally known reputation from a reputable security services provider, or at least a locally known, cached reputation for that user. In other words, it is relatively uncommon for a user to visit a website or URL for the first time.

(27) Thus, a first level of filtering may comprise simply identifying the relatively small percentage of websites that do not already have a known reputation. However, in a large enterprise, even this relatively small percentage of websites can provide a substantial workload to enterprise security scanners.

(28) Thus, additional filtering mechanisms may also be provided. For example, URLs may be filtered on contextual attributes that indicate relative likelihood of being suspicious. A first line of inquiry may be whether the website uses plain, unencrypted hypertext transfer protocol (HTTP), or HTTP Secure (HTTPS). While unencrypted HTTP was popular in the early days of the internet, in modern practice, the use of unencrypted HTTP may be considered per se suspicious, and may warrant additional analysis.

(29) Furthermore, even if the website uses HTTPS, additional attributes may be examined. For example, if the website uses HTTPS with a self-signed transport layer security (TLS) certificate, this may be deemed suspicious. If it uses HTTPS with a certificate that does not match the domain name, this may also be deemed suspicious.

(30) In another example, the URL may use HTTPS, and it may provide a valid certificate that matches the issuer, but there could be other attributes that are suspicious. For example, some web services provide "blog" or personal webpage subdomains that are hosted on a subdomain of the primary user's domain. Thus, a webpage may have a valid certificate that matches the issuer, but this may be because it is a hosted subdomain on a blog service. Thus, a listing may be maintained of known blog services that provide such subdomains, and any subdomain hosted by such a service may be blocked, or may be subject to additional scrutiny, as though it were a non-HTTPS domain.

(31) Another contextual attribute may be whether the linked webpage includes input fields, and whether any of those input fields have associated labels such as "username," "password," "credit card number," "Social Security number," or other fields for entering sensitive or personal information. A webpage that includes such fields may be selected for further analysis.

(32) One issue with this first level scanning is that phishing attackers and cyber criminals may deliberately obfuscate the source code for their webpages to hide the fact that they are providing such fields. Furthermore, where they use logos, backgrounds, colors, and/or fonts to visually mimic the appearance of reputable websites, they may also provide obfuscation so that it is difficult for traditional pattern matching scanners to identify the use of these elements.

(33) However, "digital eye" and computer vision machine learning models are not so easily defeated. Indeed, as described above, in the context of computer vision, the very visual similarity

that inspires trust in the human user is an Achilles' heel under the scrutiny of computer vision analysis. By nature, for a phishing attack to be successful, it may provide close visual similarity. Although code obfuscation may be beneficial against pattern matching scanners, visual obfuscation may need to be minimal. For example, the label on a password field may be obfuscated by using an image as the label instead of a text label. Thus, a pattern matching scanner that looks for a text field with a label similar to "password" may be defeated if a remote image is loaded, for example, from yet another third-party website. However, from a computer vision perspective, the label (however it is rendered) needs to appear to the end user to be a simple, unambiguous password label. Even a small degree of visual obfuscation may make the label appear "funny" to the end user, and may hint to the end user that something is "off" about the webpage. It is theoretically possible to create a password label that a human user can successfully decode as a password label, but that is indecipherable to a computer (indeed, this is the theory behind so-called completely automated public Turing tests to tell computers and humans apart, or CAPTCHAs). But such a visual puzzle is, as a practical matter, not useful to a phishing scammer, who is trying to give the human user as few visual clues as possible that something is wrong.

(34) For a computer vision analysis to be as successful as possible—particularly in light of code obfuscation that may make pattern matching style analysis difficult—it is beneficial to render the webpage as nearly as possible as it may be displayed to the human user. To this end, a headless browser may be used within an enterprise gateway or e-mail scanner. The headless browser includes the libraries and rendering engines provided by a traditional interactive browser. But instead of displaying the webpage to a human user for interaction, the headless browser merely renders the webpage internally. This internal rendering can then be provided to a computer vision machine learning model that has been visually trained on a number of legitimate websites. For this visual training, preference and/or precedence may be given to websites that are common targets of phishing scams. This may include, for example, banks, e-commerce sites, video streaming services, electronic payment sites, other media services, and other popular paid services.

(35) If a sophisticated phishing URL is sufficiently visually similar to, for example, the Amazon.com website to fool a human user, the computer vision system may "recognize" the website visually as the Amazon website. Once the computer vision system has visually identified the website as belonging to a certain service or provider, a scanning engine can then determine if, for example, the website is hosted on an HTTPS domain registered to and owned by that entity. If the website is visually identified as belonging to the service provider, but it is not registered to the service provider, then it can effectively be identified as malicious and can be blocked.

(36) In some embodiments, the computer vision analysis may be aided and made faster by focusing on certain elements, particularly elements that are most likely to provide visual cues to the human user. This can include the use of fonts, logos, background colors, background patterns, and other such elements that help a user intuitively identify the source of the webpage. If the rendered webpage is abstracted into these elements, then the computer vision analysis can focus on those abstracted elements rather than looking at the entire webpage as a whole.

(37) A system and method for providing visual detection of phishing websites via headless browser will now be described with more particular reference to the attached FIGURES. It should be noted that throughout the FIGURES, certain reference numerals may be repeated to indicate that a particular device or block is referenced multiple times across several FIGURES. In other cases, similar elements may be given new numbers in different FIGURES. Neither of these practices is intended to require a particular relationship between the various embodiments disclosed. In certain examples, a genus or class of elements may be referred to by a reference numeral ("widget **10**"), while individual species or examples of the element may be referred to by a hyphenated numeral ("first specific widget **10-1**" and "second specific widget **10-2**").

(38) FIG. **1** is a block diagram of a security ecosystem **100**. Embodiments of security ecosystem **100** may be configured or adapted to provide visual detection of phishing websites via headless

browser, as disclosed in the present specification.

(39) In the example of FIG. 1, security ecosystem **100** may be an enterprise, a government entity, a data center, a telecommunications provider, a “smart home” with computers, smart phones, and various internet of things (IoT) devices, or any other suitable ecosystem. Security ecosystem **100** is provided herein as an illustrative and nonlimiting example of a system that may employ, and benefit from, the teachings of the present specification.

(40) Within security ecosystem **100**, one or more users **120** operate one or more client devices **110**. A single user **120** and single client device **110** are illustrated here for simplicity, but a home or enterprise may have multiple users, each of which may have multiple devices, such as desktop computers, laptop computers, smart phones, tablets, hybrids, or similar.

(41) Client devices **110** may be communicatively coupled to one another and to other network resources via local network **170**. Local network **170** may be any suitable network or combination of one or more networks operating on one or more suitable networking protocols, including a local area network, a home network, an intranet, a virtual network, a wide area network, a wireless network, a cellular network, or the internet (optionally accessed via a proxy, virtual machine, or other similar security mechanism) by way of nonlimiting example. Local network **170** may also include one or more servers, firewalls, routers, switches, security appliances, antivirus servers, or other network devices, which may be single-purpose appliances, virtual machines, containers, or functions. Some functions may be provided on client devices **110**.

(42) In this illustration, local network **170** is shown as a single network for simplicity, but in some embodiments, local network **170** may include any number of networks, such as one or more intranets connected to the internet. Local network **170** may also provide access to an external network, such as the internet, via external network **172**. External network **172** may similarly be any suitable type of network.

(43) Local network **170** may connect to the internet via gateway **108**, which may be responsible, among other things, for providing a logical boundary between local network **170** and external network **172**. Local network **170** may also provide services such as dynamic host configuration protocol (DHCP), gateway services, router services, and switching services, and may act as a security portal across local boundary **104**.

(44) Local network **170** may also include a number of discrete IoT devices. For example, local network **170** may include IoT functionality to control lighting **132**, thermostats or other environmental controls **134**, a security system **136**, and any number of other devices **140**. Other devices **140** may include, as illustrative and nonlimiting examples, network attached storage (NAS), computers, printers, smart televisions, smart refrigerators, smart vacuum cleaners and other appliances, and network connected vehicles.

(45) Local network **170** may communicate across local boundary **104** with external network **172**. Local boundary **104** may represent a physical, logical, or other boundary. External network **172** may include, for example, websites, servers, network protocols, and other network-based services. In one example, an attacker **180** (or other similar malicious or negligent actor) also connects to external network **172**. A security services provider **190** may provide services to local network **170**, such as security software, security updates, network appliances, or similar. For example, McAfee, LLC provides a comprehensive suite of security services that may be used to protect local network **170** and the various devices connected to it.

(46) It may be a goal of users **120** to successfully operate devices on local network **170** without interference from attacker **180**. In one example, attacker **180** is a malware author whose goal or purpose is to cause malicious harm or mischief, for example, by injecting malicious object **182** into client device **110**. Once malicious object **182** gains access to client device **110**, it may try to perform work such as social engineering of user **120**, a hardware-based attack on client device **110**, modifying storage **150** (or volatile memory), modifying client application **112** (which may be running in memory), or gaining access to local resources. Furthermore, attacks may be directed at

IoT objects. IoT objects can introduce new security challenges, as they may be highly heterogeneous, and in some cases may be designed with minimal or no security considerations. To the extent that these devices have security, it may be added on as an afterthought. Thus, IoT devices may in some cases represent new attack vectors for attacker **180** to leverage against local network **170**.

(47) In some examples, attacker **180** may deliver malicious object **182** via e-mail, or by otherwise directing user **120** to the malicious website.

(48) Gateway **108** may, in some embodiments, include mechanisms to protect client devices **110** from attacker **180**. For example, gateway **108** may have a phishing website detection engine that is designed to check for elements similar to how a seasoned human user may verify a phishing website. This is beneficial, because not all users are experts in identifying phishing sites. Furthermore, cyber criminals may be very intelligent or sophisticated, and may be able to convince all but the most seasoned and experienced users to provide the requested details. The provision of a phishing website mitigation engine within gateway **108** can help to protect from this damage.

(49) The phishing website mitigation engine of gateway **108** may use a headless browser to identify a phishing website that has not yet been encountered within the enterprise. Performance may be enhanced by relying on URL reputations as a first pass check for whether a website is already known. In some embodiments, the phishing website mitigation engine of gateway **108** may verify website attributes and provide a mapping of aspects, such as a company's logo, to a suspected website. This can provide very high accuracy in identifying a phishing website.

(50) This approach realizes advantages, because targeted phishing attacks tend to be short-lived, with a time to live on the order of approximately four hours. This makes it difficult to mitigate such attacks via web crawlers, which take on the order of approximately two days to identify a new phishing website. By the time the website is identified, it may already be stale, and may have moved to a different URL.

(51) This also realizes advantages over some existing web controllers that only block risky sites. While this may be a reasonable approach for some contexts—the fact that a site has never been encountered does not per se indicate that it is malicious—it does leave users vulnerable. In particular, these solutions do not provide real-time detection of phishing websites.

(52) Enterprises may spend a great deal of money training employees, but even with training, the majority of users fail to identify actual attacks when they happen. This approach also realizes advantages over solutions that use regular expressions and pattern matching, or other machine learning methods that are error-prone in detecting phishing attacks.

(53) The phishing website mitigation engine of the present specification detects advanced zero day phishing attacks. There is no need to depend on external sources or inputs to detect a phishing site. The customer can customize a scanning process by whitelisting certain URLs and blacklisting other URLs. Detection is provided in real-time, and the accuracy rate is very high compared to some other solutions. This can, therefore, achieve high performance and is relatively less false-positive prone than some existing solutions.

(54) Malicious harm or mischief may take the form of installing root kits or other malware on client devices **110** to tamper with the system, installing spyware or adware to collect personal and commercial data, defacing websites, operating a botnet such as a spam server, or simply to annoy and harass users **120**. Thus, one aim of attacker **180** may be to install his malware on one or more client devices **110** or any of the IoT devices described. As used throughout this specification, malicious software (“malware”) includes any object configured to provide unwanted results or do unwanted work. In many cases, malware objects may be executable objects, including, by way of nonlimiting examples, viruses, Trojans, zombies, rootkits, backdoors, worms, spyware, adware, ransomware, dialers, payloads, malicious browser helper objects, tracking cookies, loggers, or similar objects designed to take a potentially-unwanted action, including, by way of nonlimiting example, data destruction, data denial, covert data collection, browser hijacking, network proxy or

redirection, covert tracking, data logging, keylogging, excessive or deliberate barriers to removal, contact harvesting, and unauthorized self-propagation. In some cases, malware could also include negligently-developed software that causes such results even without specific intent.

(55) In enterprise contexts, attacker **180** may also want to commit industrial or other espionage, such as stealing classified or proprietary data, stealing identities, or gaining unauthorized access to enterprise resources. Thus, attacker **180**'s strategy may also include trying to gain physical access to one or more client devices **110** and operating them without authorization, so that an effective security policy may also include provisions for preventing such access.

(56) In another example, a software developer may not explicitly have malicious intent, but may develop software that poses a security risk. For example, a well-known and often-exploited security flaw is the so-called buffer overrun, in which a malicious user is able to enter an overlong string into an input form and thus gain the ability to execute arbitrary instructions or operate with elevated privileges on a computing device. Buffer overruns may be the result, for example, of poor input validation or use of insecure libraries, and in many cases arise in nonobvious contexts. Thus, although not malicious, a developer contributing software to an application repository or programming an IoT device may inadvertently provide attack vectors for attacker **180**. Poorly-written applications may also cause inherent problems, such as crashes, data loss, or other undesirable behavior. Because such software may be desirable itself, it may be beneficial for developers to occasionally provide updates or patches that repair vulnerabilities as they become known. However, from a security perspective, these updates and patches are essentially new objects that must themselves be validated.

(57) Local network **170** may contract with or subscribe to a security services provider **190**, which may provide security services, updates, antivirus definitions, patches, products, and services. McAfee, LLC is a nonlimiting example of such a security services provider that offers comprehensive security and antivirus solutions. In some cases, security services provider **190** may include a threat intelligence capability such as the Global Threat Intelligence (GTI™) database provided by McAfee, LLC, or similar competing products. Security services provider **190** may update its threat intelligence database by analyzing new candidate malicious objects as they appear on client networks and characterizing them as malicious or benign.

(58) Other security considerations within security ecosystem **100** may include parents' or employers' desire to protect children or employees from undesirable content, such as pornography, adware, spyware, age-inappropriate content, advocacy for certain political, religious, or social movements, or forums for discussing illegal or dangerous activities, by way of nonlimiting example.

(59) In this example, attacker **180** operates a phishing website at the URL <http://abcbank.com.go.7rqxpt4.co.IL>. User **120** is a customer of ABC Bank, and regularly uses client devices **110** to access website **184** which provides services for ABC Bank, namely at the URL <https://www.abcbank.com>.

(60) To try to compromise personally identifying information (PII), account information, or other sensitive information from user **120**, attacker **180** may design the false website to be visually very similar to the homepage, a login page, or other portion of legitimate website **184**. Thus, attacker **180** may have as one goal to induce user **120** to visit <http://abcbank.com.go.7rqxpt4.co.IL>. For example, attacker **180** may provide the link in an e-mail, as a phishing advertisement, or by some other mechanism to try to get user **120** to click on the link. When user **120** sees the link, he may see "abcbank.com," gloss over the rest of the URL, and assume that this is a legitimate website. If user **120** is sufficiently credulous, he may attempt to log into the false website, and thus provide his username and password to attacker **180**. Furthermore, he may be tricked into providing other useful information, such as an account number, a birthday, answers to security questions, or other sensitive information.

(61) Security services operator **190** may operate a global repository or cache of URL reputations,

and may have security servers that are very capable of identifying phishing websites and blocking them. However, this URL service may require millions or even billions of URL submissions every week. After the URLs have been filtered with basic mechanisms (e.g., known/unknown, duplicated, whitelisted, blacklisted, malformed, etc.), all of the URLs with unknown reputations may be treated with the same level of priority, according to some known techniques in the art. In other words, the URL queue is a first in/first out (FIFO) queue. Under this scenario, a legitimate URL could be processed in a matter of minutes, while a phishing URL submitted moments later could be processed hours later, depending on the service load. Thus, by the time security services provider **190** analyzes the URL `http://abcbank.com.go.7rqxpt4.co.IL`, the phishing site may have already moved to a new location, and the information may be stale. This is illustrated in FIG. 2.

(62) FIG. 2 is a block diagram illustrating mitigation of an example phishing attack. In the embodiment of FIG. 2, a phishing website mitigation engine **224** includes a visual analysis engine **220**. Visual analysis engine **220** may include a “digital eye” or a computer vision style artificial intelligence (AI) model. This model may be trained on a cached copy **204** of a particular website. In practice, visual analysis engine **220** may be trained on cached copies of a large number of websites that are likely to be targets of phishing attacks. This may include, for example, online banking services, digital entertainment services, e-commerce, e-payment, and other online services that require users to provide PII or other sensitive information.

(63) To effectively deceive end users, the author of a phishing website may mimic as closely as possible the visual style of cached copy **204** of the actual website. This visual mimicry instills in the end user a sense of trust that they are accessing a legitimate website. To further instill this sense of trust, the attacker may use a deceptive URL. For example, if the attacker is mimicking `www.abcbank.com`, they may use a deceptive URL such as `www.abcbank.com.go.7rqxpt4.co.IL`. When human users see this address, they tend to focus on the portion of the URL that says “`www.abcbank.com`.” The rest of the subdomain may be subconsciously dismissed as the type of HTTP GET code that commonly appears in URLs. This HTTP GET code may provide useful information to the web browser, but provides no meaningful information to a human user. Thus, human users naturally filter out this portion of the address. Because of this natural filtering, the user does not realize that this URL is actually hosted on the domain `7rqxpt4.co.IL`. This domain is unlikely to be a legitimate domain owned and operated by ABC Bank, or its online banking subsidiary.

(64) Knowing that the website may be subjected to analysis by a corporate gateway that uses pattern matching, the attacker may also take steps to obfuscate the code for the deceptive website.

(65) Thus, remote server **206** serves a webpage at `www.abcbank.com.go.7rqxpt4.co.IL`. In its obfuscated source code form, this page may be as nearly meaningless to a pattern matching scanner as it is to a human user.

(66) Thus, to render the served page in a usable form, visual analysis engine **220** includes a headless browser **208** that renders the webpage as though it were being observed by and interacting with a human user. This provides a rendering **212**. Notably, rendering **212** is visually very similar to cached copy **204** of the legitimate webpage served by `www.abcbank.com`.

(67) In some embodiments, elements of rendering **212** may be abstracted, such as logos, images, fonts, backgrounds, and color schemes. In other embodiments, rendering **212** is presented as a unitary image.

(68) Visual analysis engine **220** of phishing website mitigation engine **224** observes rendering **212** to determine whether it can be identified as an example of a cached copy of a webpage **204**.

(69) Computer vision systems are quite effective at this kind of observation. For example, computer vision systems in recent years have been trained on images of humans, dogs, cats, fruits, or other objects. After being trained on a large data set of such images, a computer vision system is presented with a new image not in its training set. The task of the computer vision system is then to quickly classify the object. In other words, the computer vision system asks, “Is the image I am

looking at a picture of a person?” (Or, in other cases, a picture of a dog, a cat, a fruit, a car, or some other object.) A well-trained computer vision model can determine with high accuracy whether the image it is “looking” at is, in fact, an example of the trained data set.

(70) One advantage of the present approach is that rendering **212** may deliberately be designed to look very similar to cached copy **204** of a legitimate website. This visual identity, or near identity, is key to successfully deceiving the human user. If the website looks “funny” or “off,” then the human user is unlikely to be deceived.

(71) This means that one of the weaknesses of computer vision systems is of less concern in this approach. Although computer vision systems have become very good at identifying most objects belonging to a class, they are not perfect, and they are not as good as most real humans. This is the theory behind many “CAPTCHA” puzzles. If an object is obscured or distorted just the right amount, it is still recognizable to a human user (perhaps with some effort), but it is not recognizable to a computer vision system that has not been well-trained on such distorted images.

However, a distorted rendering **212** would essentially be self-defeating. If rendering **212** is distorted from cached copy **204** sufficiently to defeat a digital eye looking at the images, it is unlikely to deceive a human user-particularly, a human user with even a modicum of training.

(72) FIG. **3** is a block diagram illustrating a zero day phishing analysis pipeline **300**. Zero day phishing analysis pipeline **300** includes a number of stages of analysis for determining whether an unknown URL represents a zero day phishing attack. Note that a number of operations and/or modules are disclosed within zero day phishing analysis pipeline **300**. It is not intended for this to imply that any of these modules or operations are necessarily required in a particular embodiment. Rather, the appended claims should stand on their own in defining which elements are required for a particular embodiment.

(73) Within zero day phishing analysis pipeline **300**, a trust store **304** may be queried initially to determine whether a particular URL already has a known reputation. In some embodiments, this reputation is represented as a trust score.

(74) In decision block **308**, the system may compare a score returned by trust store **304** with a particular threshold. In some cases, the score returned by trust store **304** may indicate simply that the URL is known to be malicious (e.g., the URL has been encountered before and has been classified with a high degree of confidence as being deceptive, malicious, or otherwise harmful). In that case, control may proceed directly to terminator **390**, where the URL is detected as being malicious, and appropriate remedial action may be taken.

(75) Returning to decision block **308**, in an embodiment, another option is that the score is either above or below a particular threshold.

(76) If the score is above the trust threshold, then the website is trusted and may be treated as a known, good website. In this case, flow control may flow directly to terminator **394**, where the website is passed, and the user is permitted to interact with the website.

(77) In at least some deployments, the flow from trust store **304**, to decision block **308**, to terminator **394** may represent the vast majority of websites that a user interacts with. On a day-to-day basis, it is uncommon for a user to visit a URL for the first time. It is even more uncommon for it to be the first time the user visits that URL within a given enterprise, and yet even more uncommon that it is the first time the URL has ever been visited globally. Depending on the security infrastructure in place, there may therefore be a cached reputation for the URL indicating that the URL is known, good, and trusted. Because flow normally passes from trust store **304**, to decision block **308**, to terminator **394**, there is little disruption to the user's normal daily use of e-mail and the internet.

(78) The second most common path may be from trust store **304**, to decision block **308**, to detection terminator **390**. This represents the second most common scenario, where a website is known to the individual or the enterprise, and is known to be malicious. In this case, the known malicious website can simply be detected and blocked. Indeed, if the known malicious website is

included in a link in an e-mail, the e-mail may simply be deleted by a corporate e-mail scanner, or placed in the user's spam folder so that there is little likelihood that the user may ever even see the e-mail. Again, this represents minimal disruption of the user's daily interaction with e-mail and the internet, particularly for legitimate purposes.

(79) It should be noted that trust store **304** may be a local cache on the user's machine, an enterprise cache, and/or a global reputation store such as the GTI™ database provided by MCAFEE, LLC. In some embodiments, multiple layers of caching may be used. A local trust cache may be maintained on the user's local machine to avoid making queries to a global or enterprise trust store unnecessarily. For an enterprise, a local trust store may be maintained on an enterprise gateway that monitors all traffic in and out of the enterprise. The enterprise gateway can quickly query its local trust store for known reputations cached for specific URLs, in and out of the enterprise. In the case of a global trust store, the individual user machine or the enterprise gateway may query the global trust store only when a URL is encountered that does not have a locally or enterprise cached reputation.

(80) If control flows from decision block **308** to decision block **312**, this means that the URL does not yet have a reliable reputation. This may mean that this is the first time the URL has been accessed on a local, enterprise, and/or global level, or it may mean that the URL has been encountered before, but has not yet received sufficient analysis to have a trustworthy score that either places it above the threshold, or that marks it as a known malicious website.

(81) In block **312**, an additional filter examines the URL to determine whether it has high risk attributes. This may include observing several factors such as, by way of nonlimiting example: Does the URL use HTTP instead of HTTPS? If the website uses HTTPS, does it have a valid certificate that is assigned to this URL? Is the certificate, if any, self-signed? If the website has a valid certificate, does it belong to a subdomain of hosted personal webpages or blogs? Does the source code indicate obvious input fields for a user to input information? If there are input fields, are any of these visibly marked as being input fields for entering personal or sensitive information such as a username, password, bank account, credit card number, Social Security number, or any other sensitive or personal information?

(82) In decision block **312**, if the URL does not have high risk attributes, then it may be safe to pass the website at terminator **394**. In this context, passing the website does not necessarily mean that the website is good or desirable. It simply means that the risk of a phishing attack is low. For example, if the website is an annoying advertisement that the user would like to dismiss, this is something that the user may be able to handle on his or her own. The primary concern for purposes of this particular example is whether the URL is likely to collect sensitive personal or enterprise data from the user.

(83) Because a website without input fields cannot collect data from the user, the website can be passed even if it doesn't have a known good reputation.

(84) Thus, block **316** is encountered in the case where a URL does not have a known good score, and where the URL exhibits high risk attributes. Because this is expected to be a relatively small percentage of URLs encountered by the user or the enterprise, zero day phishing analysis pipeline **300** is expected to provide a streamlined analysis that provides minimal interruption to the user's daily personal and business activities.

(85) If the conditions of blocks **308** and **312** are met, then in block **316**, the system renders the webpage in a headless browser. Headless browser **316** includes a rendering engine similar or identical to the one that would be used by a human user. For example, headless browser **316** may include an embodiment of the WebKit web browser engine or Chrome browser core. However, headless browser **316** is not an interactive browser for users, but rather renders the webpage for analysis within zero day phishing analysis pipeline **300**.

(86) A visual analysis engine **320** observes the output of headless browser **316**. Visual analysis engine **320** may include a machine learning model that has been trained on the actual websites of a

number of web resources that may be high value targets for phishing attacks. Visual analysis engine **320** includes a digital eye that can determine whether the output of headless browser **316** “looks like” a known website that may be the target of a phishing attack. If the digital eye of headless browser **316** determines that the rendered website looks like one of the common phishing targets, then in decision block **324**, zero day phishing analysis pipeline **300** may determine that the website is malicious. If, on the other hand, the URL does not look like a website that might be a target for a phishing attack, then it can be marked as not deceptive.

(87) Thus, in decision block **324**, if the website has been marked as deceptive, then in terminator **390**, the website is detected as malicious and appropriate remedial action may be taken.

Furthermore, a reputation update may be provided to trust store **304** so that when the URL is encountered in the future, it may have a known malicious reputation.

(88) Returning to decision block **324**, if the website has been analyzed and been found not to be deceptive, then in terminator **394**, the website is passed and the user is permitted to interact with the website. Again, in some cases, a reputation update may be provided to trust store **304** so that the next time the URL is encountered, there is no need for analysis by visual analysis engine **320**.

(89) FIG. **4** is a block diagram illustrating the use of a phishing analysis pipeline **416** in an enterprise e-mail server **408**.

(90) In this example, enterprise e-mail server **408** receives an incoming e-mail **404**. As a matter of course, all incoming e-mails **404** may be scanned by enterprise e-mail server **408** to determine whether they include malicious content. In the case of enterprise e-mail server **408**, it may be more reasonable to scan all incoming e-mails, because e-mail is not expected to be a real-time communication medium. Thus, a slight latency in processing of e-mails is more acceptable to enterprise or personal users than a delay in loading webpages within a web browser.

(91) A link scanner **412** scans incoming e-mail **404** for any URLs of links provided within e-mail **404**. These links may be provided to a phishing analysis pipeline **416**. Phishing analysis pipeline **416** may, in some cases, be an embodiment of zero day phishing analysis pipeline **300** of FIG. **3**. Phishing analysis pipeline **416** may load all of the provided URLs in a headless browser and observe the loaded pages to determine whether they are apparently malicious. It should also be noted that phishing analysis pipeline **416** may have pre-analysis filters, such as a score threshold and analysis for high risk attributes.

(92) In decision block **420**, it is determined whether any of the incoming URLs were detected as malicious.

(93) If one or more URLs are detected as malicious, then in terminator **494**, the e-mail may be blocked. Depending on the configuration, the e-mail may be deleted, not delivered, or sent to a user's spam folder.

(94) Returning to decision block **420**, if the URL is not detected as malicious, then in terminator **490**, the e-mail is passed on to the user.

(95) FIG. **5** is an illustration of an enterprise gateway **508**. In this example, enterprise gateway **508** may also include a phishing analysis pipeline **516**. Again, phishing analysis pipeline **516** may be an example or embodiment of zero day phishing analysis pipeline **300** of FIG. **3**, or any other suitable phishing analysis pipeline.

(96) In this case, endpoint **504** requests a website from remote server **512**. When enterprise gateway **508** receives the request to access the URL, it may pass the URL through phishing analysis pipeline **516**. If the URL is detected as deceptive, then the request may be blocked, and optionally, a user of endpoint **504** may be notified. On the other hand, if the URL is not detected as malicious, then the user may be permitted to communicate with remote server **512**.

(97) Note that in this embodiment, the use of pre-analysis filters as illustrated in FIG. **3** may help to ensure that the user's browsing experience and normal personal and enterprise use of the internet is not interrupted by phishing analysis pipeline **516**.

(98) FIGS. **6-7** are a flowchart of a method **600** of performing zero day phishing analysis.

(99) At block **604**, the system detects an incoming e-mail with URL links, or there is a request by an endpoint for access to connect to a particular URL. This triggers the need for analysis of the URL.

(100) In block **608**, the system queries an existing reputation store **612** for a reputation for the URL. Reputation store **612** may be a local machine level URL cache, an enterprise level URL cache, a global URL cache, or any combination of the foregoing.

(101) In decision block **616**, the system determines whether the URL under analysis is known to be malicious.

(102) If the URL is known to be malicious, then following off-page connector **3** to FIG. 7, the URL is detected at terminator **690**. Upon detection, some appropriate action may be taken, such as remedial action against the known malicious URL.

(103) Returning to decision block **616** of FIG. 6, if the URL is not known malicious, then in block **620**, the system determines whether the URL has a reputation greater than a reputation threshold.

(104) If the reputation is higher than the threshold, then following off-page connector **2** to FIG. 7, at terminator **694**, the URL is passed. This means that the user is allowed to interact with the URL, and/or the e-mail is permitted to be delivered. Note that in some embodiments, this is expected to be the most common path. The URL is detected as known not malicious, and the user is permitted to interact.

(105) Returning to decision block **620** of FIG. 6, if the reputation is not greater than the threshold, then this means that either the URL is being encountered for the first time (at a personal, enterprise, and/or global level), or that the URL has been encountered before and has not yet received a reliable reputation score.

(106) In decision block **624**, the source code for the URL is analyzed to determine whether it has risky attributes, as discussed above. This could be the use of HTTP, the use of an invalid or suspicious certificate, the use of input fields, the use of labels for PII or sensitive information on input fields, or other risk attributes.

(107) The various risk attributes may be weighted appropriately to the embodiment and to particular enterprise needs. For example, the mere fact that the URL has a valid HTTPS certificate may not necessarily mean that it is to be passed. In various embodiments, the risk attributes may be assigned certain weights and combined in different ways.

(108) If the overall result of decision block **624** is that the URL is not deemed risky, then following off-page connector **2** to FIG. 7, in terminator **694** the URL or the e-mail is passed and the user is permitted to interact.

(109) Returning to decision block **624** of FIG. 6, if the URL is deemed to have sufficient risky attributes, then in block **628**, the URL is rendered in a headless browser.

(110) Following off-page connector **1** to block **632**, in some embodiments, certain user interface (UI) elements may be abstracted from the rendered webpage. These can include, by way of illustrative and nonlimiting example, a corporate logo, a background style, and the use of fonts, background colors, and other visual indicia that may be used to deceive a user into thinking that a webpage is provided by a particular provider.

(111) In block **634**, a computer vision or digital eye is applied to either the abstracted elements or to the rendered webpage as a whole.

(112) The computer vision system accesses a site store **638** that includes cached versions of webpages for high value targets of phishing attacks. Note that this may include not only the current version of a website, but also past versions going back for a particular time, such as six months, one year, five years, 10 years, or more. The number of cached versions of a website to maintain may, in some cases, depend on the value of the website. In the case of a website for a large banking venture, there may be many cached versions stored, because of the increased risk. For a website that is less popular and that collects less personal information, fewer cached versions may be maintained.

(113) In block **642**, the digital eye or computer vision system analyzes the abstracted website elements, or the rendered website as a whole, to determine whether it can visually identify the website as being an example of a known website in the site store **638**. If the website is identified as visually being an example of a known website, and if it does not have valid credentials (e.g., a URL owned by the owner of that website, and/or a digital certificate owned by the owner of that website), then the website may be detected as a phishing website.

(114) If the website is not detected, then at terminator **694**, the website is passed and the user is permitted to interact with the URL or to access the e-mail.

(115) Returning to decision block **642**, if a visual identification is made, then the website may be deemed to be a malicious phishing website, and at terminator **690**, the website is detected. Upon detection, appropriate remedial action may be taken, such as warning the user, warning an enterprise security administrator, blocking the website, blocking the e-mail, or taking other remedial action.

(116) FIG. **8** is a block diagram of a hardware platform **800**. Embodiments of hardware platform **800** may be configured or adapted to provide visual detection of phishing websites via headless browser, as disclosed in the present specification.

(117) Although a particular configuration is illustrated here, there are many different configurations of hardware platforms, and this embodiment is intended to represent the class of hardware platforms that can provide a computing device. Furthermore, the designation of this embodiment as a “hardware platform” is not intended to require that all embodiments provide all elements in hardware. Some of the elements disclosed herein may be provided, in various embodiments, as hardware, software, firmware, microcode, microcode instructions, hardware instructions, hardware or software accelerators, or similar. Furthermore, in some embodiments, entire computing devices or platforms may be virtualized, on a single device, or in a data center where virtualization may span one or a plurality of devices. For example, in a “rackscale architecture” design, disaggregated computing resources may be virtualized into a single instance of a virtual device. In that case, all of the disaggregated resources that are used to build the virtual device may be considered part of hardware platform **800**, even though they may be scattered across a data center, or even located in different data centers.

(118) Hardware platform **800** is configured to provide a computing device. In various embodiments, a “computing device” may be or comprise, by way of nonlimiting example, a computer, workstation, server, mainframe, virtual machine (whether emulated or on a “bare metal” hypervisor), network appliance, container, IoT device, high performance computing (HPC) environment, a data center, a communications service provider infrastructure (e.g., one or more portions of an Evolved Packet Core), an in-memory computing environment, a computing system of a vehicle (e.g., an automobile or airplane), an industrial control system, embedded computer, embedded controller, embedded sensor, personal digital assistant, laptop computer, cellular telephone, internet protocol telephone, smart phone, tablet computer, convertible tablet computer, computing appliance, receiver, wearable computer, handheld calculator, or any other electronic, microelectronic, or microelectromechanical device for processing and communicating data. At least some of the methods and systems disclosed in this specification may be embodied by or carried out on a computing device.

(119) In the illustrated example, hardware platform **800** is arranged in a point-to-point (PtP) configuration. This PtP configuration is popular for personal computer (PC) and server-type devices, although it is not so limited, and any other bus type may be used.

(120) Hardware platform **800** is an example of a platform that may be used to implement embodiments of the teachings of this specification. For example, instructions could be stored in storage **850**. Instructions could also be transmitted to the hardware platform in an ethereal form, such as via network interface **848**, or retrieved from another source via any suitable interconnect. Once received (from any source), the instructions may be loaded into memory **804**, and may then

be executed by one or more processor **802** to provide elements such as an operating system **806**, operational agents **808**, or data **812**.

(121) Hardware platform **800** may include several processors **802**. For simplicity and clarity, only processors PROC0 **802-1** and PROC1 **802-2** are shown. Additional processors (such as 2, 4, 8, 16, 24, 32, 64, or 128 processors) may be provided as necessary, while in other embodiments, only one processor may be provided. Details of processors **802** are not illustrated in this FIGURE, but one embodiment is illustrated in FIG. **10**. Processors may have any number of cores, such as 1, 2, 4, 8, 16, 24, 32, 64, or 128 cores.

(122) Processors **802** may be any type of processor and may communicatively couple to chipset **816** via, for example, PtP interfaces. Chipset **816** may also exchange data with other elements, such as a high performance graphics adapter **822**. In alternative embodiments, any or all of the PtP links illustrated in FIG. **8** could be implemented as any type of bus, or other configuration rather than a PtP link. In various embodiments, chipset **816** may reside on the same die or package as a central processor unit (CPU) **1112** or on one or more different dies or packages. Each chipset may support any suitable number of CPUs **802**. A chipset **816** (which may be a chipset, uncore, Northbridge, Southbridge, or other suitable logic and circuitry) may also include one or more controllers to couple other components to one or more CPUs.

(123) Two memories, **804-1** and **804-2** are shown, connected to PROC0 **802-1** and PROC1 **802-2**, respectively. As an example, each processor is shown connected to its memory in a direct memory access (DMA) configuration, though other memory architectures are possible, including ones in which memory **804** communicates with processor **810** via a bus. For example, some memories may be connected via a system bus, or in a data center, memory may be accessible in a remote DMA (RDMA) configuration.

(124) Memory **804** may include any form of volatile or nonvolatile memory including, without limitation, magnetic media (e.g., one or more tape drives), optical media, flash, random access memory (RAM), double data rate RAM (DDR RAM) nonvolatile RAM (NVRAM), static RAM (SRAM), dynamic RAM (DRAM), persistent RAM (PRAM), data-centric (DC) persistent memory (e.g., Intel® Optane/3D-crosspoint), cache, Layer 1 (L1) or Layer 2 (L2) memory, on-chip memory, registers, virtual memory region, read-only memory (ROM), flash memory, removable media, tape drive, cloud storage, or any other suitable local or remote memory component or components. Memory **804** may be used for short, medium, and/or long-term storage. Memory **804** may store any suitable data or information utilized by platform logic. In some embodiments, memory **804** may also comprise storage for instructions that may be executed by the cores of CPUs **802** or other processing elements (e.g., logic resident on chipsets **816**) to provide functionality.

(125) In certain embodiments, memory **804** may comprise a relatively low-latency volatile main memory, while storage **850** may comprise a relatively higher-latency nonvolatile memory.

However, memory **804** and storage **850** need not be physically separate devices, and in some examples may represent simply a logical separation of function (if there is any separation at all). It should also be noted that although DMA is disclosed by way of nonlimiting example, DMA is not the only protocol consistent with this specification, and that other memory architectures are available.

(126) Certain computing devices provide main memory **804** and storage **850**, for example, in a single physical memory device, and in other cases, memory **804** and/or storage **850** are functionally distributed across many physical devices. In the case of virtual machines or hypervisors, all or part of a function may be provided in the form of software or firmware running over a virtualization layer to provide the logical function, and resources such as memory, storage, and accelerators may be disaggregated (i.e., located in different physical locations across a data center). In other examples, a device such as a network interface may provide only the minimum hardware interfaces necessary to perform its logical operation, and may rely on a software driver to provide additional necessary logic. Thus, each logical block disclosed herein is broadly intended to include one or

more logic elements configured and operable for providing the disclosed logical operation of that block. As used throughout this specification, “logic elements” may include hardware, external hardware (digital, analog, or mixed-signal), software, reciprocating software, services, drivers, interfaces, components, modules, algorithms, sensors, components, firmware, hardware instructions, microcode, programmable logic, or objects that can coordinate to achieve a logical operation.

(127) Graphics adapter **822** may be configured to provide a human-readable visual output, such as a command-line interface (CLI) or graphical desktop such as Microsoft Windows, Apple OSX desktop, or a Unix/Linux X Window System-based desktop. Graphics adapter **822** may provide output in any suitable format, such as a coaxial output, composite video, component video, video graphics array (VGA), or digital outputs such as digital visual interface (DVI), FPDLink, DisplayPort, or high definition multimedia interface (HDMI), by way of nonlimiting example. In some examples, graphics adapter **822** may include a hardware graphics card, which may have its own memory and its own graphics processing unit (GPU).

(128) Chipset **816** may be in communication with a bus **828** via an interface circuit. Bus **828** may have one or more devices that communicate over it, such as a bus bridge **832**, I/O devices **835**, network interface **848**, accelerators **846**, communication devices **840**, and a keyboard and/or mouse **838**, by way of nonlimiting example. In general terms, the elements of hardware platform **800** may be coupled together in any suitable manner. For example, a bus may couple any of the components together. A bus may include any known interconnect, such as a multi-drop bus, a mesh interconnect, a fabric, a ring interconnect, a round-robin protocol, a PtP interconnect, a serial interconnect, a parallel bus, a coherent (e.g., cache coherent) bus, a layered protocol architecture, a differential bus, or a Gunning transceiver logic (GTL) bus, by way of illustrative and nonlimiting example.

(129) Communication devices **840** can broadly include any communication not covered by network interface **848** and the various I/O devices described herein. This may include, for example, various universal serial bus (USB), FireWire, Lightning, or other serial or parallel devices that provide communications.

(130) I/O Devices **835** may be configured to interface with any auxiliary device that connects to hardware platform **800** but that is not necessarily a part of the core architecture of hardware platform **800**. A peripheral may be operable to provide extended functionality to hardware platform **800**, and may or may not be wholly dependent on hardware platform **800**. In some cases, a peripheral may be a computing device in its own right. Peripherals may include input and output devices such as displays, terminals, printers, keyboards, mice, modems, data ports (e.g., serial, parallel, USB, Firewire, or similar), network controllers, optical media, external storage, sensors, transducers, actuators, controllers, data acquisition buses, cameras, microphones, speakers, or external storage, by way of nonlimiting example.

(131) In one example, audio I/O **842** may provide an interface for audible sounds, and may include in some examples a hardware sound card. Sound output may be provided in analog (such as a 3.5 mm stereo jack), component (“RCA”) stereo, or in a digital audio format such as S/PDIF, AES3, AES47, HDMI, USB, Bluetooth, or Wi-Fi audio, by way of nonlimiting example. Audio input may also be provided via similar interfaces, in an analog or digital form.

(132) Bus bridge **832** may be in communication with other devices such as a keyboard/mouse **838** (or other input devices such as a touch screen, trackball, etc.), communication devices **840** (such as modems, network interface devices, peripheral interfaces such as PCI or PCIe, or other types of communication devices that may communicate through a network), audio I/O devices **842**, a data storage device **844**, and/or accelerators **846**. In alternative embodiments, any portions of the bus architectures could be implemented with one or more PtP links.

(133) Operating system **806** may be, for example, Microsoft Windows, Linux, UNIX, Mac OS X, IOS, MS-DOS, or an embedded or real-time operating system (including embedded or real-time

flavors of the foregoing). In some embodiments, a hardware platform **800** may function as a host platform for one or more guest systems that invoke application (e.g., operational agents **808**).

(134) Operational agents **808** may include one or more computing engines that may include one or more non-transitory computer-readable mediums having stored thereon executable instructions operable to instruct a processor to provide operational functions. At an appropriate time, such as upon booting hardware platform **800** or upon a command from operating system **806** or a user or security administrator, processor **802** may retrieve a copy of the operational agent (or software portions thereof) from storage **850** and load it into memory **804**. Processor **810** may then iteratively execute the instructions of operational agents **808** to provide the desired methods or functions.

(135) As used throughout this specification, an “engine” includes any combination of one or more logic elements, of similar or dissimilar species, operable for and configured to perform one or more methods provided by the engine. In some cases, the engine may be or include a special integrated circuit designed to carry out a method or a part thereof, a field-programmable gate array (FPGA) programmed to provide a function, a special hardware or microcode instruction, other programmable logic, and/or software instructions operable to instruct a processor to perform the method. In some cases, the engine may run as a “daemon” process, background process, terminate-and-stay-resident program, a service, system extension, control panel, bootup procedure, basic input/output system (BIOS) subroutine, or any similar program that operates with or without direct user interaction. In certain embodiments, some engines may run with elevated privileges in a “driver space” associated with ring 0, 1, or 2 in a protection ring architecture. The engine may also include other hardware, software, and/or data, including configuration files, registry entries, application programming interfaces (APIs), and interactive or user-mode software by way of nonlimiting example.

(136) Where elements of an engine are embodied in software, computer program instructions may be implemented in programming languages, such as an object code, an assembly language, or a high-level language such as OpenCL, FORTRAN, C, C++, JAVA, or HTML. These may be used with any compatible operating systems or operating environments. Hardware elements may be designed manually, or with a hardware description language such as Spice, Verilog, and VHDL. The source code may define and use various data structures and communication messages. The source code may be in a computer executable form (e.g., via an interpreter), or the source code may be converted (e.g., via a translator, assembler, or compiler) into a computer executable form, or converted to an intermediate form such as byte code. Where appropriate, any of the foregoing may be used to build or describe appropriate discrete or integrated circuits, whether sequential, combinatorial, state machines, or otherwise.

(137) Network interface **848** may be provided to communicatively couple hardware platform **800** to a wired or wireless network or fabric. A “network,” as used throughout this specification, may include any communicative platform operable to exchange data or information within or between computing devices, including, by way of nonlimiting example, a local network, a switching fabric, an ad-hoc local network, Ethernet (e.g., as defined by the IEEE 802.3 standard), Fibre Channel, InfiniBand, Wi-Fi, or other suitable standard. Intel® Omni-Path™ Architecture (OPA), TrueScale™, Ultra Path Interconnect (UPI) (formerly called QPI or KTI), FibreChannel, Ethernet, FibreChannel over Ethernet (FCOE), InfiniBand, PCI, PCIe, fiber optics, millimeter wave guide, an internet architecture, a packet data network (PDN) offering a communications interface or exchange between any two nodes in a system, a local area network (LAN), metropolitan area network (MAN), wide area network (WAN), wireless local area network (WLAN), virtual private network (VPN), intranet, plain old telephone system (POTS), or any other appropriate architecture or system that facilitates communications in a network or telephonic environment, either with or without human interaction or intervention. Network interface **848** may include one or more physical ports that may couple to a cable (e.g., an Ethernet cable, other cable, or waveguide).

(138) In some cases, some or all of the components of hardware platform **800** may be virtualized,

in particular the processor(s) and memory. For example, a virtualized environment may run on OS **806**, or OS **806** could be replaced with a hypervisor or virtual machine manager. In this configuration, a virtual machine running on hardware platform **800** may virtualize workloads. A virtual machine in this configuration may perform essentially all of the functions of a physical hardware platform.

(139) In a general sense, any suitably-configured processor can execute any type of instructions associated with the data to achieve the operations illustrated in this specification. Any of the processors or cores disclosed herein could transform an element or an article (for example, data) from one state or thing to another state or thing. In another example, some activities outlined herein may be implemented with fixed logic or programmable logic (for example, software and/or computer instructions executed by a processor).

(140) Various components of the system depicted in FIG. **8** may be combined in a system-on-a-chip (SoC) architecture or in any other suitable configuration. For example, embodiments disclosed herein can be incorporated into systems including mobile devices such as smart cellular telephones, tablet computers, personal digital assistants, portable gaming devices, and similar. These mobile devices may be provided with SoC architectures in at least some embodiments. An example of such an embodiment is provided in FIG. **9**. Such an SoC (and any other hardware platform disclosed herein) may include analog, digital, and/or mixed-signal, radio frequency (RF), or similar processing elements. Other embodiments may include a multichip module (MCM), with a plurality of chips located within a single electronic package and configured to interact closely with each other through the electronic package. In various other embodiments, the computing functionalities disclosed herein may be implemented in one or more silicon cores in application-specific integrated circuits (ASICs), FPGAs, and other semiconductor chips.

(141) FIG. **9** is a block illustrating selected elements of an example SoC **900**. Embodiments of SoC **900** may be configured or adapted to provide visual detection of phishing websites via headless browser, as disclosed in the present specification.

(142) At least some of the teachings of the present specification may be embodied on an SoC **900**, or may be paired with an SoC **900**. SoC **900** may include, or may be paired with, an advanced reduced instruction set computer machine (ARM) component. For example, SoC **900** may include or be paired with any ARM core, such as A-9, A-15, or similar. This architecture represents a hardware platform that may be useful in devices such as tablets and smartphones, by way of illustrative example, including Android phones or tablets, iPhone (of any version), iPad, Google Nexus, Microsoft Surface. SoC **900** could also be integrated into, for example, a PC, server, video processing components, laptop computer, notebook computer, netbook, or touch-enabled device.

(143) As with hardware platform **800** above, SoC **900** may include multiple cores **902a** and **902b**. In this illustrative example, SoC **900** also includes an L2 cache control **904**, a GPU **906**, a video codec **908**, a liquid crystal display (LCD) I/F **910** and an interconnect **912**. L2 cache control **904** can include a bus interface unit **914**, a L2 cache **916**. Liquid crystal display (LCD) I/F **910** may be associated with mobile industry processor interface (MIPI)/HDMI links that couple to an LCD.

(144) SoC **900** may also include a subscriber identity module (SIM) I/F **918**, a boot ROM **920**, a synchronous dynamic random access memory (SDRAM) controller **922**, a flash controller **924**, a serial peripheral interface (SPI) master **928**, a suitable power control **930**, a dynamic RAM (DRAM) **932**, and flash **934**. In addition, one or more embodiments include one or more communication capabilities, interfaces, and features such as instances of Bluetooth™ **936**, a 3G modem **938**, a global positioning system (GPS) **940**, and an 802.11 Wi-Fi **942**.

(145) Designers of integrated circuits such as SoC **900** (or other integrated circuits) may use intellectual property (IP) blocks to simplify system design. An IP block is a modular, self-contained hardware block that can be easily integrated into the design. Because the IP block is modular and self-contained, the integrated circuit (IC) designer need only “drop in” the IP block to use the functionality of the IP block. The system designer can then make the appropriate connections to

inputs and outputs.

(146) IP blocks are often “black boxes.” In other words, the system integrator using the IP block may not know, and need not know, the specific implementation details of the IP block. Indeed, IP blocks may be provided as proprietary third-party units, with no insight into the design of the IP block by the system integrator.

(147) For example, a system integrator designing an SoC for a smart phone may use IP blocks in addition to the processor core, such as a memory controller, a nonvolatile memory (NVM) controller, Wi-Fi, Bluetooth, GPS, a fourth or fifth-generation network (4G or 5G), an audio processor, a video processor, an image processor, a graphics engine, a GPU engine, a security controller, and many other IP blocks. In many cases, each of these IP blocks has its own embedded microcontroller.

(148) FIG. **10** is a block diagram illustrating selected elements of a processor **1000**. Embodiments of processor **1000** may be configured or adapted to provide visual detection of phishing websites via headless browser, as disclosed in the present specification.

(149) In various examples, and throughout this specification and the appended claims, a “processor” may include any combination of logic elements operable to execute instructions, whether loaded from memory, or implemented directly in hardware, including, by way of nonlimiting example, a microprocessor, microcontroller, CPU, advanced RISC (reduced instruction set computing) machine (ARM), digital signal processor (DSP), FPGA, GPU, programmable logic array, ASIC, or virtual machine processor. In certain architectures, a multi-core processor may be provided, having for example, 2, 4, 8, 12, 16, 24, 32, 64, or 128 cores. In some embodiments, one or more co-processors or accelerators (hardware or software) may also be provided for specialized or support functions. In general, processor **1000** may include any number of processing elements, which may be symmetrical or asymmetrical.

(150) Examples of hardware processing elements include: a thread unit, a thread slot, a thread, a process unit, a context, a context unit, a logical processor, a hardware thread, a core, and/or any other element, which is capable of holding a state for a processor, such as an execution state or architectural state. In other words, a processing element, in one embodiment, refers to any hardware capable of being independently associated with code, such as a software thread, operating system, application, or other code. A physical processor (or processor socket) typically refers to an integrated circuit, which potentially includes any number of other processing elements, such as cores or hardware threads.

(151) A core may refer to logic located on an integrated circuit capable of maintaining an independent architectural state, wherein each independently maintained architectural state is associated with at least some dedicated execution resources. A hardware thread may refer to any logic located on an integrated circuit capable of maintaining an independent architectural state, wherein the independently maintained architectural states share access to execution resources. A physical CPU may include any suitable number of cores. In various embodiments, cores may include one or more out-of-order processor cores or one or more in-order processor cores. However, cores may be individually selected from any type of core, such as a native core, a software managed core, a core adapted to execute a native instruction set architecture (ISA), a core adapted to execute a translated ISA, a co-designed core, or other known core. In a heterogeneous core environment (i.e. asymmetric cores), some form of translation, such as binary translation, may be utilized to schedule or execute code on one or both cores.

(152) Processor **1000** includes one or more processor cores **1002**, including core **1002-1-1002-N**. Cores **1002** may be, as appropriate, single-thread cores or multi-thread cores. In multithreaded cores, more than one hardware thread may be provided at a time, and the core may therefore provide more than one logical core per physical core. The cores may be configured to execute instruction code. Each processor **1000** may include at least one shared cache **1030**, which may be treated logically as part of memory **1040**. Caches **1030** may be filled according to known caching

techniques, and may store instructions and/or data that may be used by one or more components of processor **1000**.

(153) Processor **1000** may include an integrated memory controller (MC) **1034**, to communicate with memory **1040**. Memory controller **1034** may include logic and circuitry to interface with memory **1040**, and may also include a cache controller to handle filling and evicting instructions and data to and from cache **1030**.

(154) By way of example, each core **1002** may include front-end logic **1006**, execution logic **1014**, and backend logic **1018**.

(155) In the illustrated embodiment, front-end logic **1006** includes an instruction decoder or decoders **1008**, register renaming logic **1010**, and scheduling logic **1012**. Decoder **1008** may decode instructions received. Register renaming logic **1010** may provide register renaming, for example to facilitate pipelining. Scheduling logic **1012** may schedule instruction execution, and may provide out-of-order (OOO) execution. Front-end logic **1006** may fetch incoming instructions, perform various processing (e.g., caching, decoding, branch predicting, etc.), and pass instructions to execution logic **1014**.

(156) Execution logic **1014** includes one or more execution units **1016-1-1016-N**. Execution units **1016** may include hardware instructions and microcode to carry out the provided instructions.

(157) Backend logic **1018** includes retirement logic **1020**. Core **1002** may provide for speculative execution of instructions, branch prediction, and similar. Retirement logic **1020** may be configured to determine which predicted instructions were actually needed by the program flow.

(158) Processor **1000** may also include a PtP controller **1032**, which enables connection to an uncore, chipset, Northbridge, Southbridge, or bus, by way of example.

(159) FIG. **11** is a block diagram of a network function virtualization (NFV) infrastructure **1100**. Embodiments of NFV infrastructure **1100** may be configured or adapted to provide visual detection of phishing websites via headless browser, as disclosed in the present specification.

(160) NFV is an aspect of network virtualization that is generally considered distinct from, but that can still interoperate with, software defined networking (SDN). For example, virtual network functions (VNFs) may operate within the data plane of an SDN deployment. NFV was originally envisioned as a method for providing reduced capital expenditure (Capex) and operating expenses (Opex) for telecommunication services. One feature of NFV is replacing proprietary, special-purpose hardware appliances with virtual appliances running on commercial off-the-shelf (COTS) hardware within a virtualized environment. In addition to Capex and Opex savings, NFV provides a more agile and adaptable network. As network loads change, VNFs can be provisioned (“spun up”) or removed (“spun down”) to meet network demands. For example, in times of high load, more load balancing VNFs may be spun up to distribute traffic to more workload servers (which may themselves be virtual machines). In times when more suspicious traffic is experienced, additional firewalls or deep packet inspection (DPI) appliances may be needed.

(161) Because NFV started out as a telecommunications feature, many NFV instances are focused on telecommunications. However, NFV is not limited to telecommunication services. In a broad sense, NFV includes one or more VNFs running within a network function virtualization infrastructure (NFVI), such as NFVI **400**. Often, the VNFs are inline service functions that are separate from workload servers or other nodes. These VNFs can be chained together into a service chain, which may be defined by a virtual subnetwork, and which may include a serial string of network services that provide behind-the-scenes work, such as security, logging, billing, and similar.

(162) In the example of FIG. **11**, an NFV orchestrator **1101** manages a number of the VNFs **1112** running on an NFVI **1100**. NFV requires nontrivial resource management, such as allocating a very large pool of compute resources among appropriate numbers of instances of each VNF, managing connections between VNFs, determining how many instances of each VNF to allocate, and managing memory, storage, and network connections. This may require complex software

management, thus making NFV orchestrator **1101** a valuable system resource. Note that NFV orchestrator **1101** may provide a browser-based or graphical configuration interface, and in some embodiments may be integrated with SDN orchestration functions.

(163) Note that NFV orchestrator **1101** itself may be virtualized (rather than a special-purpose hardware appliance). NFV orchestrator **1101** may be integrated within an existing SDN system, wherein an operations support system (OSS) manages the SDN. This may interact with cloud resource management systems (e.g., OpenStack) to provide NFV orchestration. An NFVI **1100** may include the hardware, software, and other infrastructure to enable VNFs to run. This may include a hardware platform **1102** on which one or more VMs **1104** may run. For example, hardware platform **1102-1** in this example runs VMs **1104-1** and **1104-2**. Hardware platform **1102-2** runs VMs **1104-3** and **1104-4**. Each hardware platform may include a hypervisor **1120**, virtual machine manager (VMM), or similar function, which may include and run on a native (bare metal) operating system, which may be minimal so as to consume very few resources.

(164) Hardware platforms **1102** may be or comprise a rack or several racks of blade or slot servers (including, e.g., processors, memory, and storage), one or more data centers, other hardware resources distributed across one or more geographic locations, hardware switches, or network interfaces. An NFVI **1100** may also include the software architecture that enables hypervisors to run and be managed by NFV orchestrator **1101**.

(165) Running on NFVI **1100** are a number of VMS **1104**, each of which in this example is a VNF providing a virtual service appliance. Each VM **1104** in this example includes an instance of the Data Plane Development Kit (DPDK), a virtual operating system **1108**, and an application providing the VNF **1112**.

(166) Virtualized network functions could include, as nonlimiting and illustrative examples, firewalls, intrusion detection systems, load balancers, routers, session border controllers, DPI services, network address translation (NAT) modules, or call security association.

(167) The illustration of FIG. **11** shows that a number of VNFs **1104** have been provisioned and exist within NFVI **1100**. This FIGURE does not necessarily illustrate any relationship between the VNFs and the larger network, or the packet flows that NFVI **1100** may employ.

(168) The illustrated DPDK instances **1116** provide a set of highly-optimized libraries for communicating across a virtual switch (vSwitch) **1122**. Like VMs **1104**, vSwitch **1122** is provisioned and allocated by a hypervisor **1120**. The hypervisor uses a network interface to connect the hardware platform to the data center fabric (e.g., an HFI). This HFI may be shared by all VMs **1104** running on a hardware platform **1102**. Thus, a vSwitch may be allocated to switch traffic between VMs **1104**. The vSwitch may be a pure software vSwitch (e.g., a shared memory vSwitch), which may be optimized so that data are not moved between memory locations, but rather, the data may stay in one place, and pointers may be passed between VMs **1104** to simulate data moving between ingress and egress ports of the vSwitch. The vSwitch may also include a hardware driver (e.g., a hardware network interface IP block that switches traffic, but that connects to virtual ports rather than physical ports). In this illustration, a distributed vSwitch **1122** is illustrated, wherein vSwitch **1122** is shared between two or more physical hardware platforms **1102**.

(169) The foregoing outlines features of several embodiments so that those skilled in the art may better understand various aspects of the present disclosure. The embodiments disclosed can readily be used as the basis for designing or modifying other processes and structures to carry out the teachings of the present specification. Any equivalent constructions to those disclosed do not depart from the spirit and scope of the present disclosure. Design considerations may result in substitute arrangements, design choices, device possibilities, hardware configurations, software implementations, and equipment options.

(170) In certain embodiments, some of the components illustrated herein may be omitted or consolidated. In a general sense, the arrangements depicted in the FIGURES may be more logical

in their representations, whereas a physical architecture may include various permutations, combinations, and/or hybrids of these elements.

(171) With the numerous examples provided herein, interaction may be described in terms of two, three, four, or more electrical components. These descriptions are provided for purposes of clarity and example only. Any of the illustrated components, modules, and elements of the FIGURES may be combined in various configurations, all of which fall within the scope of this specification.

(172) In certain cases, it may be easier to describe one or more functionalities by disclosing only selected element. Such elements are selected to illustrate specific information to facilitate the description. The inclusion of an element in the FIGURES is not intended to imply that the element must appear in the invention, as claimed, and the exclusion of certain elements from the FIGURES is not intended to imply that the element is to be excluded from the invention as claimed. Similarly, any methods or flows illustrated herein are provided by way of illustration only. Inclusion or exclusion of operations in such methods or flows should be understood the same as inclusion or exclusion of other elements as described in this paragraph. Where operations are illustrated in a particular order, the order is a nonlimiting example only. Unless expressly specified, the order of operations may be altered to suit a particular embodiment.

(173) Other changes, substitutions, variations, alterations, and modifications will be apparent to those skilled in the art. All such changes, substitutions, variations, alterations, and modifications fall within the scope of this specification.

(174) In order to aid the United States Patent and Trademark Office (USPTO) and, any readers of any patent or publication flowing from this specification, the Applicant: (a) does not intend any of the appended claims to invoke paragraph (f) of 35 U.S.C. section 112, or its equivalent, as it exists on the date of the filing hereof unless the words “means for” or “steps for” are specifically used in the particular claims; and (b) does not intend, by any statement in the specification, to limit this disclosure in any way that is not otherwise expressly reflected in the appended claims, as originally presented or as amended.

EXAMPLE IMPLEMENTATIONS

(175) There is disclosed, in one example, a computing apparatus, comprising: a processor and a memory; a network interface; and instructions encoded within the memory to instruct the processor to: receive a uniform resource locator (URL) for analysis, the URL to access a web page via a remote server; via the network interface, retrieve from the remote server a copy of the web page; render the web page in a headless browser to provide a computer-accessible visual output; perform visual analysis of the visual output via a digital eye; compare the visual analysis to a plurality of known phishing target websites; and if the comparison identifies the web page as visually similar to a known phishing target website, detect the web page as a phishing web page.

(176) There is further disclosed an example computing apparatus, wherein the instructions are further to provide a pre-filter stage before retrieving the web page from the remote server.

(177) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises checking whether the URL has an existing cached local, enterprise, or global reputation.

(178) There is further disclosed an example computing apparatus, wherein the instructions are to pass the URL if it has a known benign reputation.

(179) There is further disclosed an example computing apparatus, wherein the instructions are to block the URL if it has a known malicious or phishing reputation.

(180) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis only if it has input fields.

(181) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if the input fields collect potentially private or sensitive information.

(182) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if the purpose of the input fields cannot be

determined.

(183) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if it includes non-encrypted hypertext transfer protocol (HTTP).

(184) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page provides a self-signed transport layer security (TLS) certificate.

(185) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page provides a transport layer security (TLS) certificate that does not match the URL.

(186) There is further disclosed an example computing apparatus, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page is hosted on a personal blog or personal webpage subdomain of a reputable URL.

(187) There is further disclosed an example computing apparatus, wherein performing visual analysis comprises abstracting out a selected element of the visual output.

(188) There is further disclosed an example computing apparatus, wherein the selected element is a logo.

(189) There is further disclosed an example computing apparatus, wherein the selected element is a font.

(190) There is further disclosed an example computing apparatus, wherein the selected element is a background color or image.

(191) There is further disclosed an example computing apparatus, wherein the instructions are part of a web gateway.

(192) There is further disclosed an example computing apparatus, wherein the instructions are part of an e-mail scanner.

(193) There is also disclosed an example of one or more tangible, non-transitory computer-readable media having stored thereon executable instructions to: select a uniform resource locator (URL) for visual analysis, the URL to access a web page under analysis via a remote server; retrieve a copy of the web page under analysis from the remote server; visually analyze the web page under analysis, comprising: visually rendering the web page under analysis in a headless browser; applying a computer vision machine learning model to the visual rendering, the machine learning model trained on a plurality of web pages known to be phishing targets; within the computer vision machine learning model, detecting the web page under analysis as an instance of a web page known to be a phishing target; detecting the web page under analysis as a phishing attack; and blocking the web page under analysis.

(194) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein selecting the URL for visual analysis comprises applying a pre-filter stage.

(195) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises checking whether the URL has an existing cached local, enterprise, or global reputation.

(196) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the instructions are to pass the URL if it has a known benign reputation.

(197) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the instructions are to block the URL if it has a known malicious or phishing reputation.

(198) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis only if it has input fields.

(199) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if the

input fields collect potentially private or sensitive information.

(200) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if the purpose of the input fields cannot be determined.

(201) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if it includes non-encrypted hypertext transfer protocol (HTTP).

(202) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page provides a self-signed transport layer security (TLS) certificate.

(203) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page provides a transport layer security (TLS) certificate that does not match the URL.

(204) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the pre-filter stage comprises identifying the URL for visual analysis if the web page is hosted on a personal blog or personal webpage subdomain of a reputable URL.

(205) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein performing visual analysis comprises abstracting out a selected element of the visual output.

(206) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the selected element is a logo.

(207) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the selected element is a font.

(208) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the selected element is a background color or image.

(209) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the instructions are part of a web gateway.

(210) There is further disclosed an example of one or more tangible, non-transitory computer-readable media, wherein the instructions are part of an e-mail scanner.

(211) There is also disclosed an example computer-implemented method of providing zero-day-capable phishing detection for unknown URLs, comprising: receiving an incoming URL providing a reference to a web page; determining that the URL does not have a known reliable reputation; applying a pre-filtering stage to the URL; upon determining that the URL did not pass the pre-filtering stage, subjecting the URL to visual machine learning analysis, comprising: rendering the web page in a headless browser to provide a machine-readable visual representation of the web page; instructing a machine learning model to identify the web page, wherein the machine learning model has been trained on a set of web page captures of online services known to be phishing targets; receiving an identification from the machine learning model; and if the machine learning model has identified the web page as being identical or substantially similar to a web service known to be a phishing target, detecting the URL as a phishing attack, and assigning the URL a malicious reputation.

(212) There is further disclosed an example method, wherein the pre-filter stage comprises identifying the URL for visual analysis if it has input fields.

(213) There is further disclosed an example method, wherein the pre-filter stage comprises identifying the URL for visual analysis if the input fields collect potentially private or sensitive information.

(214) There is further disclosed an example method, wherein the pre-filter stage comprises identifying the URL for visual analysis if the purpose of the input fields cannot be determined.

(215) There is further disclosed an example method, wherein the pre-filter stage comprises identifying the URL for visual analysis if it includes non-encrypted hypertext transfer protocol

(HTTP).

(216) There is further disclosed an example method, wherein the pre-filter stage comprises determining that the web page provides a self-signed transport layer security (TLS) certificate.

(217) There is further disclosed an example method, wherein the pre-filter stage comprises determining that the URL provides a transport layer security (TLS) certificate that does not match the URL.

(218) There is further disclosed an example method, wherein the pre-filter stage comprises determining that the web page is hosted on a personal blog or personal webpage subdomain of a reputable URL.

(219) There is further disclosed an example method, wherein performing visual analysis comprises abstracting out a selected element of the visual output.

(220) There is further disclosed an example method, wherein the selected element is a logo.

(221) There is further disclosed an example method, wherein the selected element is a font.

(222) There is further disclosed an example method, wherein the selected element is a background color or image.

(223) There is further disclosed an example method, further comprising providing the method as part of a web gateway.

(224) There is further disclosed an example method, further comprising providing the method as part of an e-mail scanner.

(225) There is further disclosed an example apparatus comprising means for performing the method of a number of the above examples.

(226) There is further disclosed an example apparatus, wherein the means for performing the method comprise a processor and a memory.

(227) There is further disclosed an example apparatus, wherein the memory comprises machine-readable instructions, that when executed cause the apparatus to perform the method of a number of the above examples.

(228) There is further disclosed an example apparatus, wherein the apparatus is a computing system.

(229) There is further disclosed an example of at least one computer-readable medium comprising instructions that, when executed, implement a method or realize an apparatus as illustrated in a number of the above examples.

(230) There is further disclosed an computer-implemented method of protecting a user from phishing attacks, comprising: detecting that the user has opened a website associated with a uniform resource locator (URL), and determining that the URL does not have a known good reputation; rendering a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspecting the test version of the website with a digital eye, and determining that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warning the user that the website is or may be a phishing website.

(231) There is further disclosed an example, further comprising querying a cloud reputation service to determine that the URL does not have a known good reputation.

(232) There is further disclosed an example, further comprising querying a local reputation cache to determine that the URL does not have a known good reputation.

(233) There is further disclosed an example, further comprising determining that the user has not previously visited the URL to determine that the URL does not have a known good reputation.

(234) There is further disclosed an example, wherein abstracting the at least one visual element includes a corporate logo.

(235) There is further disclosed an example, wherein abstracting the at least one visual element includes a font.

(236) There is further disclosed an example, wherein abstracting the at least one visual element

includes a background color.

(237) There is further disclosed an example, wherein abstracting the at least one visual element includes a background pattern.

(238) There is further disclosed an example, wherein abstracting the at least one visual element includes a color scheme.

(239) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL is a non-secure hypertext transport protocol (HTTP) URL.

(240) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL uses transport layer security (TLS), but that a certificate of the URL is not associated with the known legitimate website.

(241) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL belongs to a blogs or personal webpages subdomain of a known third-party service not associated with the known legitimate website.

(242) There is further disclosed an example of an apparatus comprising means for performing the method.

(243) There is further disclosed an example, wherein the means for performing the method comprise a processor and a memory.

(244) There is further disclosed an example, wherein the memory comprises machine-readable instructions that, when executed, cause the apparatus to perform the method.

(245) There is further disclosed an example, wherein the apparatus is a computing system.

(246) There is further disclosed an example of at least one computer readable medium comprising instructions that, when executed, implement a method or realize an apparatus described.

(247) There is further disclosed an example of one or more tangible, nontransitory computer-readable storage media having stored thereon executable instructions to: detect that a user has opened a website associated with a uniform resource locator (URL), and determine that the URL does not have a known good reputation; render a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspect the test version of the website with a digital eye, and determine that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warn the user that the website is or may be a phishing website.

(248) There is further disclosed an example, wherein the instructions are further to query a cloud reputation service to determine that the URL does not have a known good reputation.

(249) There is further disclosed an example, wherein the instructions are further to query a local reputation cache to determine that the URL does not have a known good reputation.

(250) There is further disclosed an example, wherein the instructions are further to determine that the user has not previously visited the URL to determine that the URL does not have a known good reputation.

(251) There is further disclosed an example, wherein abstracting the at least one visual element includes a corporate logo.

(252) There is further disclosed an example, wherein abstracting the at least one visual element includes a font.

(253) There is further disclosed an example, wherein abstracting the at least one visual element includes a background color.

(254) There is further disclosed an example, wherein abstracting the at least one visual element includes a background pattern.

(255) There is further disclosed an example, wherein abstracting the at least one visual element includes a color scheme.

(256) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL is a non-secure hypertext transport protocol (HTTP) URL.

(257) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL uses transport layer security (TLS), but that a certificate of the URL is not associated with the known legitimate website.

(258) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL belongs to a blogs or personal webpages subdomain of a known third-party service not associated with the known legitimate website.

(259) There is further disclosed an example of a computing apparatus, comprising: a hardware platform comprising a processor circuit and a memory; and instructions encoded within the memory to instruct the processor circuit to: detect that a user has opened a website associated with a uniform resource locator (URL), and determine that the URL does not have a known good reputation; render a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspect the test version of the website with a digital eye, and determine that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warn the user that the website is or may be a phishing website.

(260) There is further disclosed an example, wherein the instructions are further to query a cloud reputation service to determine that the URL does not have a known good reputation.

(261) There is further disclosed an example, wherein the instructions are further to query a local reputation cache to determine that the URL does not have a known good reputation.

(262) There is further disclosed an example, wherein the instructions are further to determine that the user has not previously visited the URL to determine that the URL does not have a known good reputation.

(263) There is further disclosed an example, wherein abstracting the at least one visual element includes a corporate logo.

(264) There is further disclosed an example, wherein abstracting the at least one visual element includes a font.

(265) There is further disclosed an example, wherein abstracting the at least one visual element includes a background color.

(266) There is further disclosed an example, wherein abstracting the at least one visual element includes a background pattern.

(267) There is further disclosed an example, wherein abstracting the at least one visual element includes a color scheme.

(268) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL is a non-secure hypertext transport protocol (HTTP) URL.

(269) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL uses transport layer security (TLS), but that a certificate of the URL is not associated with the known legitimate website.

(270) There is further disclosed an example, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL belongs to a blogs or personal webpages subdomain of a known third-party service not associated with the known legitimate website.

(271) There is further disclosed an example, wherein the computing apparatus is a desktop computer.

- (272) There is further disclosed an example, wherein the computing apparatus is a workstation.
- (273) There is further disclosed an example, wherein the computing apparatus is a laptop computer.
- (274) There is further disclosed an example, wherein the computing apparatus is a notebook computer.
- (275) There is further disclosed an example, wherein the computing apparatus is a netbook.
- (276) There is further disclosed an example, wherein the computing apparatus is a tablet computer.
- (277) There is further disclosed an example, wherein the computing apparatus is a convertible tablet computer.
- (278) There is further disclosed an example, wherein the computing apparatus is a smart phone.
- (279) There is further disclosed an example, wherein the computing apparatus is an Android phone.
- (280) There is further disclosed an example, wherein the computing apparatus is an iPhone.
- (281) There is further disclosed an example, wherein the computing apparatus is a Windows phone.
- (282) There is further disclosed an example, wherein the computing apparatus is a server.
- (283) There is further disclosed an example, further comprising a guest infrastructure to realize server functions.
- (284) There is further disclosed an example, wherein the guest infrastructure comprises virtualization infrastructure.
- (285) There is further disclosed an example, wherein the guest infrastructure comprises containerization infrastructure.
- (286) There is further disclosed an example, wherein the computing apparatus is a gateway.

Claims

1. A computer-implemented method of protecting a user from phishing attacks, comprising: detecting that the user has opened a website associated with a uniform resource locator (URL), and determining that the URL does not have a known good reputation, comprising querying a cloud-based reputation store for the URL; in response to the determining, rendering a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspecting the test version of the website with a digital eye, the digital eye comprising a machine learning model visually trained on known legitimate websites, and determining that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warning the user that the website is or may be a phishing website.
2. The method of claim 1, further comprising querying a cloud reputation service to determine that the URL does not have a known good reputation.
3. The method of claim 1, further comprising querying a local reputation cache to determine that the URL does not have a known good reputation.
4. The method of claim 1, further comprising determining that the user has not previously visited the URL to determine that the URL does not have a known good reputation.
5. The method of claim 1, wherein abstracting the at least one visual element includes a corporate logo.
6. The method of claim 1, wherein abstracting the at least one visual element includes a font.
7. The method of claim 1, wherein abstracting the at least one visual element includes a background color.
8. The method of claim 1, wherein abstracting the at least one visual element includes a background pattern.
9. The method of claim 1, wherein abstracting the at least one visual element includes a color scheme.
10. The method of claim 1, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL is a

non-secure hypertext transport protocol (HTTP) URL.

11. The method of claim 1, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL uses transport layer security (TLS), but that a certificate of the URL is not associated with the known legitimate website.

12. The method of claim 1, wherein determining that the test version of the website looks like the known legitimate website not associated with the URL comprises determining that the URL belongs to a blogs or personal webpages subdomain of a known third-party service not associated with the known legitimate website.

13. One or more tangible, nontransitory computer-readable storage media having stored thereon executable instructions to: detect that a user has opened a website associated with a uniform resource locator (URL), and determine that the URL does not have a known good reputation, comprising querying a cloud-based reputation store for the URL; in response to the determining, render a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspect the test version of the website with a digital eye, the digital eye comprising a machine learning model visually trained on known legitimate websites, and determine that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warn the user that the website is or may be a phishing website.

14. The one or more tangible, nontransitory computer-readable storage media of claim 13, wherein the instructions are further to query a cloud reputation service to determine that the URL does not have a known good reputation.

15. The one or more tangible, nontransitory computer-readable storage media of claim 13, wherein the instructions are further to query a local reputation cache to determine that the URL does not have a known good reputation.

16. The one or more tangible, nontransitory computer-readable storage media of claim 13, wherein the instructions are further to determine that the user has not previously visited the URL to determine that the URL does not have a known good reputation.

17. A computing apparatus, comprising: a hardware platform comprising a processor circuit and a memory; and instructions encoded within the memory to instruct the processor circuit to: detect that a user has opened a website associated with a uniform resource locator (URL), and determine that the URL does not have a known good reputation, comprising querying a cloud-based reputation store for the URL; in response to the determining, render a test version of the website in a headless web browser, including abstracting the website into at least one visual element of the test version; visually inspect the test version of the website with a digital eye, the digital eye comprising a machine learning model visually trained on known legitimate websites, and determine that the test version of the web looks like a known legitimate website not associated with the URL; and based on the visual inspection, warn the user that the website is or may be a phishing website.

18. The computing apparatus of claim 17, wherein the instructions are further to query a cloud reputation service to determine that the URL does not have a known good reputation.

19. The computing apparatus of claim 17, wherein the instructions are further to query a local reputation cache to determine that the URL does not have a known good reputation.

20. The computing apparatus of claim 17, wherein the computing apparatus is a gateway.
