

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent
Kind Code
Date of Patent
Inventor(s)

12413722
B2
September 09, 2025
Guleryuz; Onur et al.

Filtering with side-information using contextually-designed filters

Abstract

Obtaining a restored frame from a degraded frame includes obtaining, for a pixel of the degraded frame, magnitude features based on a first window centered at the pixel. A cardinality N of the magnitude features is at least 1. The magnitude features are used to obtain a pixel-adaptive filter. The pixel-adaptive filter is applied to the pixel to obtain a pixel of the restored frame.

Inventors: Guleryuz; Onur (San Francisco, CA), Mukherjee; Debargha (Cupertino, CA)
Applicant: Google LLC (Mountain View, CA)
Family ID: 83191979
Assignee: GOOGLE LLC (Mountain View, CA)
Appl. No.: 18/686155
Filed (or PCT Filed): August 12, 2022
PCT No.: PCT/US2022/040232
PCT Pub. No.: WO2023/043552
PCT Pub. Date: March 23, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240357098 A1	Oct. 24, 2024

Related U.S. Application Data

us-provisional-application US 63244867 20210916

Publication Classification

Int. Cl.: H04N7/12 (20060101); H04N19/117 (20140101); H04N19/124 (20140101); H04N19/136 (20140101); H04N19/82 (20140101)

U.S. Cl.:

CPC **H04N19/117** (20141101); **H04N19/124** (20141101); **H04N19/136** (20141101); **H04N19/82** (20141101);

Field of Classification Search

CPC: H04N (19/182)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2011/0090242	12/2010	Cote et al.	N/A	N/A
2018/0131968	12/2017	Mukherjee	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
3516870	12/2018	EP	N/A

OTHER PUBLICATIONS

International Search Report and Written Opinion of International Application No. PCT/US2022/040232, dated Nov. 10, 2022, 13 pgs. cited by applicant

Primary Examiner: Lee; Y

Attorney, Agent or Firm: Young Basile Hanlon & MacFarlane, P.C.

Background/Summary

BACKGROUND

(1) Digital video streams can represent video using a sequence of frames or still images. Digital video can be used for various applications including, for example, video conferencing, high-definition video entertainment, video advertisements, or sharing of user-generated videos. A digital video stream can contain a large amount of data and consume a significant amount of computing or communication resources of a computing device for processing, transmission, or storage of the video data. Various approaches have been proposed to reduce the amount of data in video streams, including compression and other encoding techniques.

(2) Encoding using compression can be performed by breaking frames or images into blocks that are then compressed, often using encoding techniques that result in loss of some data. A decoder can apply one or more filters to a reconstructed frame to remove or smooth out artifacts caused by (e.g., lossy) encoding.

SUMMARY

(3) The disclosure relates in general to video coding, and in particular to filtering with side-information using contextually-designed filters.

(4) A first aspect is a method for obtaining a restored frame from a degraded frame. The method includes obtaining, for a pixel of the degraded frame, magnitude features based on a first window centered at the pixel. A cardinality N of the magnitude features is at least 1. The method further includes using the magnitude features to obtain a pixel-adaptive filter and applying the pixel-adaptive filter to the pixel to obtain a pixel of the restored frame.

(5) A second aspect is a method that includes obtaining a filter bank including filters, obtaining, for pixels of a degraded frame, respective sets of combining scalars for combining the filters of the filter bank, obtaining, for the pixels of the degraded frame, respective pixel-specific filters by combining the filters of

the filter bank using the respective sets of combining scalars, and obtaining a restored frame by filtering the pixels of the degraded frame using the respective pixel-specific filters.

(6) A third aspect is a device that includes a processor. The processor is configured to obtain, for a pixel of a degraded frame, magnitude features based on a first window centered at the pixel, use the magnitude features to obtain a pixel-adaptive filter, and apply the pixel-adaptive filter to the pixel to obtain a pixel of a restored frame.

(7) These and other aspects of the present disclosure are disclosed in the following detailed description of the embodiments, the appended claims, and the accompanying figures.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The description herein makes reference to the accompanying drawings wherein like reference numerals refer to like parts throughout the several views.
- (2) FIG. 1 is a schematic of a video encoding and decoding system.
- (3) FIG. 2 is a block diagram of an example of a computing device that can implement a transmitting station or a receiving station.
- (4) FIG. 3 is a diagram of a video stream to be encoded and subsequently decoded.
- (5) FIG. 4 is a block diagram of an encoder.
- (6) FIG. 5 is a block diagram of a decoder.
- (7) FIG. 6 is a flowchart of an example of a technique of filtering using combining scalars.
- (8) FIG. 7 is a flowchart of an example of obtaining a filter bank G.
- (9) FIG. 8 is a flowchart of a technique for restoring a degraded frame.
- (10) FIG. 9 is a flowchart of an example of a technique 900 for filtering with pixel-specific filters obtained using a look-up-table.
- (11) FIG. 10 illustrates an example of obtaining magnitude feature values.
- (12) FIG. 11 is an illustration of a portion of a lookup table.
- (13) FIG. 12 is a flowchart of a technique for restoring a degraded frame.

DETAILED DESCRIPTION

(14) As mentioned above, compression schemes related to coding video streams can include breaking images into blocks and generating a digital video output bitstream using one or more techniques to limit the information included in the output. A received bitstream can be decoded to re-create the blocks and the source images from the limited information. Encoding a video stream, or a portion thereof, such as a frame or a block, can include using temporal or spatial similarities in the video stream to improve coding efficiency. For example, a current block of a video stream can be encoded based on identifying a difference (residual) between previously coded pixel values and those in the current block. In this way, only the residual and/or parameters used to generate the residual need be added to the bitstream instead of including the entirety of the current block. The residual can be encoded using a lossy quantization step. Decoding (i.e., reconstructing) an encoded block from such a residual often results in a distortion between the original (i.e., source) block and the reconstructed block.

(15) Post-reconstruction loop filters can be used in various ways to improve reconstructed frames distorted or degraded as a result of the encoding and decoding processes. For example, in-loop deblocking filters can be used to modify pixel values near borders between blocks to limit the visibility of those borders within the reconstructed frame. Other loop filters can be used to bring the reconstructed images closer to the source images by, for example, adding offsets that are determined at the encoder to pixel values of the reconstructed frame. Those loop filters operate in a blind setting (i.e., without access to, or influence from, both a source frame and its associated reconstructed frame).

(16) In traditional implementations, a set of fixed filters may be available at the decoder for applying to a decoded or a reconstructed frame (collectively, a degraded frame). One or more of the available filters may be applied by the decoder. In some other traditional implementations, the decoder may receive indications (e.g., indexes) of the one or more filters that the decoder is to apply. However, such traditional implementations do not further adapt the filter weights (also referred to as taps) or combine the filters in a way that is best adapted to the frame itself. That is, while the fixed filters may be designed to generally provide average improvements over a large set of frames, such fixed filters may not and cannot take into

account peculiarities of certain frames.

(17) Implementations according to this disclosure can filter a decoded frame of a video or an image (referred to herein as a “degraded frame”) using pixel-adaptive filters to obtain a restored frame. A pixel-adaptive (or pixel-specific) filter can be obtained for at least some (e.g., each) pixel of the degraded image. The pixel-specific filter can be obtained by combining a set of filters (i.e., a filter bank) using pixel-specific combining scalars. The combining scalars for a pixel (or, equivalently, the pixel-adaptive filter for that pixel) are obtained based on local information (e.g., pixel values) in a neighborhood of the pixel (e.g., a window) that includes the pixel itself; hence the term “pixel-adaptive” filter.

(18) In an example, a filter bank can include first filters (referred to herein as “fixed filters”) available at the decoder, second filters (referred to herein as “side filters”) received from an encoder in a compressed bitstream that includes the frame, or both (i.e., fixed filters and side filters). Pixels of the degraded frame are filtered using the respective pixel-specific filters to obtain the restored image.

(19) Said another way, the combining (or combined) scalars for a pixel (more specifically, a pixel location) are obtained using the information (e.g., pixel values) at that pixel location and at least some of its surrounding (i.e., neighboring) pixels. The combining scalars can be obtained in any number of ways.

(20) In an example, the combining scalars can be obtained using a machine learning (ML) model (e.g., a neural network) that is trained to receive a degraded frame and output, in an example, the combining scalars. Side-information may be used as further described herein. In other examples, the combining scalars can be obtained using (e.g., simple) features that do not require the computation complexity of neural networks. In an example, Wiener filters can be used to obtain the features. The Wiener filters aim to increase quality especially over directional features and textures in the decoded picture.

(21) With respect to the ML model, the ML model can be trained to minimize errors between restored images and their corresponding source (i.e., original) images. To restate, the ML model can derive (e.g., calculate, infer, output, etc.) a vector of combining scalars at each pixel of the degraded frame. The respective combining scalars at each pixel are used to combine filters of a filter-bank to obtain a pixel-adaptive filter. As such, potentially different filters can be obtained at each pixel of the image. Said another way, the filter used for one pixel is independently derived from the filter derived for another pixel.

(22) Pixel-adaptive filters can be applied to the pixels of the image at a neighborhood of the pixel to arrive at the filtered value for that pixel. In an example, the combining scalars may be obtained for a restoration unit or block. A restoration unit or a restoration block can be a luma block of size 256×256 pixels or a chroma block of size 128×128 pixels. Other restoration unit sizes are possible. A restoration unit may be defined as a portion of a reconstructed frame to which an in-loop filter is to be applied.

(23) As used herein, a decoded frame or image is referred to as a “degraded frame” because it is not as close to the original (i.e., source) image as the restored frame. In an example, a decoder can filter the degraded frame with the aid of (e.g. using) side-information received in a compressed bitstream. In an example, the decoder can filter the degraded frame to obtain the restored frame in-loop (i.e., within the video compression loop). As such, the filtered pixels of the restored frame can be used in prediction of other pixels of other video frames.

(24) Implementations according to this disclosure can realize (e.g., obtain) a very large number of filters with minimal computations and side-information that results in improved rate-distortion performance at lower computational complexity. Pixel-adaptive filtering can increase quality in decoded frames. In some situations (e.g., low bit rate situations), pixel-adaptive filtering can improve performance without side-information by relying solely on finely characterized pixel contexts.

(25) As such, described herein are in-loop filtering techniques that may be used to augment or replace loop restoration processes used by codecs. Improved performance at both high and low bitrates can be obtained. The filters obtained according to this disclosure are non-separable filters. Separable filters may perform well for horizontal and vertical lines or edges. However, when a restoration unit includes directional lines (i.e., non-horizontal or non-vertical lines), separable filters do not perform well.

(26) Filtering with side-information using contextually-designed filters is described herein first with reference to a system in which the teachings can be incorporated. As alluded to above, in the restoration herein, the frame can be restored in one or more portions. Each of these portions is referred to herein respectively as a “restoration unit,” where restoration units may overlap or may not overlap each other.

(27) FIG. 1 is a schematic of a video encoding and decoding system 100. A transmitting station 102 can be, for example, a computer having an internal configuration of hardware such as that described in FIG. 2.

However, other suitable implementations of the transmitting station **102** are possible. For example, the processing of the transmitting station **102** can be distributed among multiple devices.

(28) A network **104** can connect the transmitting station **102** and a receiving station **106** for encoding and decoding of the video stream. Specifically, the video stream can be encoded in the transmitting station **102** and the encoded video stream can be decoded in the receiving station **106**. The network **104** can be, for example, the Internet. The network **104** can also be a local area network (LAN), wide area network (WAN), virtual private network (VPN), cellular telephone network or any other means of transferring the video stream from the transmitting station **102** to, in this example, the receiving station **106**.

(29) The receiving station **106**, in one example, can be a computer having an internal configuration of hardware such as that described in FIG. 2. However, other suitable implementations of the receiving station **106** are possible. For example, the processing of the receiving station **106** can be distributed among multiple devices.

(30) Other implementations of the video encoding and decoding system **100** are possible. For example, an implementation can omit the network **104**. In another implementation, a video stream can be encoded and then stored for transmission at a later time to the receiving station **106** or any other device having memory. In one implementation, the receiving station **106** receives (e.g., via the network **104**, a computer bus, and/or some communication pathway) the encoded video stream and stores the video stream for later decoding. In an example implementation, a real-time transport protocol (RTP) is used for transmission of the encoded video over the network **104**. In another implementation, a transport protocol other than RTP can be used, e.g., an HTTP-based video streaming protocol.

(31) When used in a video conferencing system, for example, the transmitting station **102** and/or the receiving station **106** can include the ability to both encode and decode a video stream as described below. For example, the receiving station **106** could be a video conference participant who receives an encoded video bitstream from a video conference server (e.g., the transmitting station **102**) to decode and view and further encodes and transmits its own video bitstream to the video conference server for decoding and viewing by other participants.

(32) FIG. 2 is a block diagram of an example of a computing device **200** that can implement a transmitting station or a receiving station. For example, the computing device **200** can implement one or both of the transmitting station **102** and the receiving station **106** of FIG. 1. The computing device **200** can be in the form of a computing system including multiple computing devices, or in the form of a single computing device, for example, a mobile phone, a tablet computer, a laptop computer, a notebook computer, a desktop computer, and the like.

(33) A CPU **202** in the computing device **200** can be a central processing unit. Alternatively, the CPU **202** can be any other type of device, or multiple devices, capable of manipulating or processing information now-existing or hereafter developed. Although the disclosed implementations can be practiced with a single processor as shown, e.g., the CPU **202**, advantages in speed and efficiency can be achieved using more than one processor.

(34) A memory **204** in the computing device **200** can be a read-only memory (ROM) device or a random-access memory (RAM) device in an implementation. Any other suitable type of storage device can be used as the memory **204**. The memory **204** can include code and data **206** that is accessed by the CPU **202** using a bus **212**. The memory **204** can further include an operating system **208** and application programs **210**, the application programs **210** including at least one program that permits the CPU **202** to perform the methods described here. For example, the application programs **210** can include applications 1 through N, which further include a video coding application that performs the methods described here. The computing device **200** can also include a secondary storage **214**, which can, for example, be a memory card used with a computing device **200** that is mobile. Because the video communication sessions can contain a significant amount of information, they can be stored in whole or in part in the secondary storage **214** and loaded into the memory **204** as needed for processing.

(35) The computing device **200** can also include one or more output devices, such as a display **218**. The display **218** can be, in one example, a touch sensitive display that combines a display with a touch sensitive element that is operable to sense touch inputs. The display **218** can be coupled to the CPU **202** via the bus **212**. Other output devices that permit a user to program or otherwise use the computing device **200** can be provided in addition to or as an alternative to the display **218**. When the output device is or includes a display, the display can be implemented in various ways, including by a liquid crystal display (LCD), a

cathode-ray tube (CRT) display or light emitting diode (LED) display, such as an organic LED (OLED) display.

(36) The computing device **200** can also include or be in communication with an image-sensing device **220**, for example a camera, or any other image-sensing device **220** now existing or hereafter developed that can sense an image such as the image of a user operating the computing device **200**. The image-sensing device **220** can be positioned such that it is directed toward the user operating the computing device **200**. In an example, the position and optical axis of the image-sensing device **220** can be configured such that the field of vision includes an area that is directly adjacent to the display **218** and from which the display **218** is visible.

(37) The computing device **200** can also include or be in communication with a sound-sensing device **222**, for example a microphone, or any other sound-sensing device now existing or hereafter developed that can sense sounds near the computing device **200**. The sound-sensing device **222** can be positioned such that it is directed toward the user operating the computing device **200** and can be configured to receive sounds, for example, speech or other utterances, made by the user while the user operates the computing device **200**.

(38) Although FIG. 2 depicts the CPU **202** and the memory **204** of the computing device **200** as being integrated into a single unit, other configurations can be utilized. The operations of the CPU **202** can be distributed across multiple machines (each machine having one or more of processors) that can be coupled directly or across a local area or other network. The memory **204** can be distributed across multiple machines such as a network-based memory or memory in multiple machines performing the operations of the computing device **200**. Although depicted here as a single bus, the bus **212** of the computing device **200** can be composed of multiple buses. Further, the secondary storage **214** can be directly coupled to the other components of the computing device **200** or can be accessed via a network and can comprise a single integrated unit such as a memory card or multiple units such as multiple memory cards. The computing device **200** can thus be implemented in a wide variety of configurations.

(39) FIG. 3 is a diagram of an example of a video stream **300** to be encoded and subsequently decoded. The video stream **300** includes a video sequence **302**. At the next level, the video sequence **302** includes a number of adjacent frames **304**. While three frames are depicted as the adjacent frames **304**, the video sequence **302** can include any number of adjacent frames **304**. The adjacent frames **304** can then be further subdivided into individual frames, e.g., a frame **306**. At the next level, the frame **306** can be divided into a series of segments **308** or planes. The segments **308** can be subsets of frames that permit parallel processing, for example. The segments **308** can also be subsets of frames that can separate the video data into separate colors. For example, the frame **306** of color video data can include a luminance plane and two chrominance planes. The segments **308** can be sampled at different resolutions.

(40) Whether or not the frame **306** is divided into the segments **308**, the frame **306** can be further subdivided into blocks **310**, which can contain data corresponding to, for example, 16×16 pixels in the frame **306**. The blocks **310** can also be arranged to include data from one or more segments **308** of pixel data. The blocks **310** can also be of any other suitable size such as 4×4 pixels, 8×8 pixels, 16×8 pixels, 8×16 pixels, 16×16 pixels or larger.

(41) FIG. 4 is a block diagram of an encoder **400** in accordance with implementations of this disclosure. The encoder **400** can be implemented, as described above, in the transmitting station **102** such as by providing a computer software program stored in memory, for example, the memory **204**. The computer software program can include machine instructions that, when executed by a processor such as the CPU **202**, cause the transmitting station **102** to encode video data in the manner described herein. The encoder **400** can also be implemented as specialized hardware included in, for example, the transmitting station **102**. The encoder **400** has the following stages to perform the various functions in a forward path (shown by the solid connection lines) to produce an encoded or compressed bitstream **420** using the video stream **300** as input: an intra/inter prediction stage **402**, a transform stage **404**, a quantization stage **406**, and an entropy encoding stage **408**. The encoder **400** can also include a reconstruction path (shown by the dotted connection lines) to reconstruct a frame for encoding of future blocks. In FIG. 4, the encoder **400** has the following stages to perform the various functions in the reconstruction path: a dequantization stage **410**, an inverse transform stage **412**, a reconstruction stage **414**, and a loop filtering stage **416**. Other structural variations of the encoder **400** can be used to encode the video stream **300**.

(42) When the video stream **300** is presented for encoding, the frame **306** can be processed in units of blocks. At the intra/inter prediction stage **402**, a block can be encoded using intra-frame prediction (also

called intra-prediction) or inter-frame prediction (also called inter-prediction), or a combination of both. In any case, a prediction block can be formed. In the case of intra-prediction, all or a part of a prediction block can be formed from samples in the current frame that have been previously encoded and reconstructed. In the case of inter-prediction, all or part of a prediction block can be formed from samples in one or more previously constructed reference frames determined using motion vectors.

(43) Next, still referring to FIG. 4, the prediction block can be subtracted from the current block at the intra/inter prediction stage **402** to produce a residual block (also called a residual). The transform stage **404** transforms the residual into transform coefficients in, for example, the frequency domain using block-based transforms. Such block-based transforms include, for example, the Discrete Cosine Transform (DCT) and the Asymmetric Discrete Sine Transform (ADST). Other block-based transforms are possible. Further, combinations of different transforms can be applied to a single residual. In one example of application of a transform, the DCT transforms the residual block into the frequency domain where the transform coefficient values are based on spatial frequency. The lowest frequency (DC) coefficient at the top-left of the matrix and the highest frequency coefficient at the bottom-right of the matrix. It is worth noting that the size of a prediction block, and hence the resulting residual block, can be different from the size of the transform block. For example, the prediction block can be split into smaller blocks to which separate transforms are applied.

(44) The quantization stage **406** converts the transform coefficients into discrete quantum values, which are referred to as quantized transform coefficients, using a quantizer value or a quantization level. For example, the transform coefficients can be divided by the quantizer value and truncated. The quantized transform coefficients are then entropy encoded by the entropy encoding stage **408**. Entropy coding can be performed using any number of techniques, including token and binary trees. The entropy-encoded coefficients, together with other information used to decode the block, which can include for example the type of prediction used, transform type, motion vectors and quantizer value, are then output to the compressed bitstream **420**. The information to decode the block can be entropy coded into block, frame, slice and/or section headers within the compressed bitstream **420**. The compressed bitstream **420** can also be referred to as an encoded video stream or encoded video bitstream, and the terms will be used interchangeably herein.

(45) The reconstruction path in FIG. 4 (shown by the dotted connection lines) can be used to ensure that both the encoder **400** and a decoder **500** (described below) use the same reference frames and blocks to decode the compressed bitstream **420**. The reconstruction path performs functions that are similar to functions that take place during the decoding process that are discussed in more detail below, including dequantizing the quantized transform coefficients at the dequantization stage **410** and inverse transforming the dequantized transform coefficients at the inverse transform stage **412** to produce a derivative residual block (also called a derivative residual). At the reconstruction stage **414**, the prediction block that was predicted at the intra/inter prediction stage **402** can be added to the derivative residual to create a reconstructed block. The loop filtering stage **416** can be applied to the reconstructed block to reduce distortion such as blocking artifacts.

(46) Other variations of the encoder **400** can be used to encode the compressed bitstream **420**. For example, a non-transform based encoder **400** can quantize the residual signal directly without the transform stage **404** for certain blocks or frames. In another implementation, an encoder **400** can have the quantization stage **406** and the dequantization stage **410** combined into a single stage.

(47) FIG. 5 is a block diagram of a decoder **500** in accordance with implementations of this disclosure. The decoder **500** can be implemented in the receiving station **106**, for example, by providing a computer software program stored in the memory **204**. The computer software program can include machine instructions that, when executed by a processor such as the CPU **202**, cause the receiving station **106** to decode video data in the manner described in FIG. 10 below. The decoder **500** can also be implemented in hardware included in, for example, the transmitting station **102** or the receiving station **106**.

(48) The decoder **500**, similar to the reconstruction path of the encoder **400** discussed above, includes in one example the following stages to perform various functions to produce an output video stream **516** from the compressed bitstream **420**: an entropy decoding stage **502**, a dequantization stage **504**, an inverse transform stage **506**, an intra/inter-prediction stage **508**, a reconstruction stage **510**, a loop filtering stage **512**, and an optional post filtering stage **514**. Other structural variations of the decoder **500** can be used to decode the compressed bitstream **420**. The loop filtering stage **512** can include a deblocking filtering stage.

(49) When the compressed bitstream **420** is presented for decoding, the data elements within the

compressed bitstream **420** can be decoded by the entropy decoding stage **502** to produce a set of quantized transform coefficients. The dequantization stage **504** dequantizes the quantized transform coefficients (e.g., by multiplying the quantized transform coefficients by the quantizer value), and the inverse transform stage **506** inverse transforms the dequantized transform coefficients using the selected transform type to produce a derivative residual that can be identical to that created by the inverse transform stage **412** in the encoder **400**. Using header information decoded from the compressed bitstream **420**, the decoder **500** can use the intra/inter-prediction stage **508** to create the same prediction block as was created in the encoder **400**, e.g., at the intra/inter prediction stage **402**. At the reconstruction stage **510**, the prediction block can be added to the derivative residual to create a reconstructed block. The loop filtering stage **512** can be applied to the reconstructed block to reduce blocking artifacts. Other filtering can be applied to the reconstructed block. In an example, the deblocking filtering stage **514** is applied to the reconstructed block to reduce blocking distortion as described below, and the result is output as an output video stream **516**. The output video stream **516** can also be referred to as a decoded video stream, and the terms will be used interchangeably herein.

(50) Other variations of the decoder **500** can be used to decode the compressed bitstream **420**. For example, the decoder **500** can produce the output video stream **516** without the post filtering stage **514**. In some implementations of the decoder **500**, the post filtering stage **514** is applied before the loop filtering stage **512**. Additionally, or alternatively, the encoder **400** includes a deblocking filtering stage in addition to the loop filtering stage **416**.

(51) FIG. **6** is a flowchart of an example of a technique **600** of filtering using combining scalars. Filtering using combining scalars can be implemented by a loop filtering stage, such as the loop filtering stage **416** of FIG. **4** or the loop filtering stage **512** of FIG. **5**. Filtering using combining scalars can be implemented by a reconstruction stage, such as the reconstruction stage **510** of FIG. **5** or the reconstruction stage **414** of FIG. **4**.

(52) A degraded frame **602** is input to an ML model **604**. In different implementations, the ML model **604** can be trained to use different types of data as inputs. In an example, the ML model **604** may be trained to use, as input, a frame that is the output of a traditional loop filtering stage (i.e., a loop filtering stage that does not include obtaining a restored frame as described herein). As such, the degraded frame **602** can be the output of the loop filtering stage. In an example, the ML model **604** may be trained to use, as input, a frame that is the output of the reconstruction stage. As such, the degraded frame **602** can be the output of the reconstruction phase. The training of the ML model **604** is further described below. The degraded frame has a size of $P \times Q$ and each pixel can be at a location (p, q) , where $p=0, \dots, P-1$ and $q=0, \dots, Q-1$. In an example, the inputs to the ML model may be restoration units. In yet another example, the ML model may partition an input frame into restoration units.

(53) The ML model **604** outputs combining scalars **606** (i.e., combining scalars $c_{sup.p,q}$) for each pixel of the degraded frame. For each pixel (p, q) (i.e., the pixel at the Cartesian location (p, q)), the ML model **604** outputs N values, where N is the number of filters of a filter bank G . Obtaining the filter bank G is described below with respect to FIG. **7**. The combining scalars $c_{sup.p,q}$ can be conveniently denoted as a column vector of size $N \times 1$. The combining vectors $c_{sup.p,q}$ for the pixel at Cartesian location (p, q) of the degraded frame can be given by equation (1). As such, the ML model **604** can output a total of $P \times Q$ combining scalars where each of the $P \times Q$ combining scalars includes N scalars (e.g., numbers, values, multipliers). As such, the ML model **604** can output a total of $P \times Q \times N$ values.

$$(54) \quad c^{p,q} = \begin{bmatrix} c_1^{p,q} \\ c_2^{p,q} \\ \vdots \\ c_N^{p,q} \end{bmatrix} \quad (1)$$

(55) At **608**, for each pixel at (p, q) , a respective filter $f_{sup.p,q}$ is obtained by combining the N filters of a filter bank **610** (i.e., the filter bank G) using the combining scalars $c_{sup.p,q}$. Each of the filters of the filter bank G includes K taps (e.g., weights)—one tap for each of the pixels of a neighborhood of the pixel (p, q) . To illustrate, if the neighborhood is a window of size $K=A \times B$ in size (e.g., $K=3 \times 5$), where A is the number of rows and B is the number of columns, then each filter of the filter bank G includes $K=A \times B$ (e.g., $K=3 \times 5=15$) weights. The window can be a 3×3 , 5×5 , 7×7 , or of some other size. The size of the window can be chosen based on a desired complexity. That is, K can be complexity dependent. While a pixel

neighborhood may be 2-dimensional, each filter (g.sub.i) of the filter bank can be represented as a 1-dimensional $K \times 1$ column vector as shown in equation (2).

$$(56) \quad g_i = \begin{bmatrix} g_{i,1} \\ g_{i,2} \\ \vdots \\ g_{i,K} \end{bmatrix} \quad (2)$$

(57) In equation (2), each weight g.sub.i,j corresponds (or is used as a multiplier for) a corresponding pixel of the window. To illustrate, assume that the window is a square window of size 3×3 . As such, $K=9$. With respect to the pixel (p, q), the window includes, in raster scan order, the pixels at Cartesian locations (p-1,q-1), (p-1,q), (p-1, q+1), (p, q-1), (p, q), (p, q+1), (p+1, q-1), (p+1, q), and (p+1, q+1). Thus, the weight g.sub.i,1 corresponds to (e.g., is used for) the pixel at location, (p-1,q-1); the weight g.sub.i,2 corresponds to the pixel at location, (p-1,q); . . . ; the weight g.sub.i,6 corresponds to pixel at location, (p, q+1); . . . ; and the weight g.sub.i,K corresponds to the pixel at location, (p+1, q+1).

(58) It is noted that, while described herein that a neighborhood of a pixel is a square or a rectangular set of pixels, the disclosure is not so limited and the neighborhood of a pixel at (p, q) can be any set of pixels that are proximal to the pixel at (p, q). For example, the neighborhood can include the pixel and its immediate (i.e., top, left, right, and bottom) neighboring pixels and not include any diagonally adjacent pixels of the pixel. In another example, the neighborhood can include pixels in a same row or pixels in a same column as the pixel, but not both. The weights of the column vector can be arranged in a lexicographical order, such as a raster scan order.

(59) FIG. 7 is a flowchart 700 of an example of obtaining a filter bank G. As mentioned above, the filter bank can include at least one of fixed filters (G.sub.fixed) or filters received in a compressed bitstream (G.sub.side).

(60) At 702, the side filters G.sub.side are optionally obtained (e.g., decoded) from a compressed bitstream, such as the compressed bitstream 420 of FIG. 5. An encoder, such as the encoder 400 of FIG. 4 may transmit the side filters G.sub.side in the compressed bitstream. The compressed bitstream may or may not include side filters. Whereas the fixed filters G.sub.fixed may generally improve most images, the side filters G.sub.side are determined by the encoder to improve the particular image (or video sequence) being decoded.

(61) The encoder may generate side filters G.sub.side in situations where the encoder determines that the compression performance can be improved using side filters G.sub.side, or more generally, filter-related information. That is, the generated side filters G.sub.side can improve the compression performance over only using (if at all) the fixed filters G.sub.fixed.

(62) Obtaining the side filters G.sub.side from the compressed bitstream can mean decoding the respective weights of the side filters G.sub.side from the compressed bitstream. As such, and assuming that the number of side filters is s, then the compressed bitstream can include up to $s \times K$ weights. The compressed bitstream can include a syntax element indicating the number of side filters to be decoded. The side filters G.sub.side can be conveniently represented, as shown in equation (3), as a matrix where each element g.sub.i is a column vector that includes the weights of the filter, as described with respect to equation (2).

$$(63) \quad G_{\text{side}} = [g_1, \text{.Math.}, g_s], \text{ where } s \text{ is the number of side filters} \quad (3)$$

(64) At 704, the fixed filters G.sub.fixed are optionally obtained (e.g., retrieved). The fixed filters G.sub.fixed can be filters that are designed a priori and can be known and/or available for use by an encoder and a decoder. The fixed filters G.sub.fixed can be designed to generally improve many different images. The fixed filters G.sub.fixed includes (N-s) filters. The fixed filters G.sub.fixed can be conveniently represented, as shown in equation (4), as a matrix where each element g.sub.i is a column vector that includes the weights of the filter, as described with respect to equation (2).

$$(65) \quad G_{\text{fixed}} = [g_{s+1}, \text{.Math.}, g_N], \text{ where } N \text{ is the total number of filters} \quad (4)$$

(66) In an example, whether to use or not to use any of the fixed filters G.sub.fixed may be transmitted in the compressed bitstream. For example, the compressed bitstream can include one or more syntax elements that the decoder can use to determine whether to use the fixed filters G.sub.fixed. In another example, the combining scalars determine whether, which, and to what extent the fixed filters G.sub.fixed are used. That is,

the values of the combining scalars determine whether, and to what extent, each fixed filter is utilized. For example, if the ML model **604** infers that the fixed filters G.sub.fixed do not improve the degraded frame **602**, then the ML model **604** may output combining scalars corresponding to the fixed filters G.sub.fixed having zero values. To illustrate, assume that the filter g.sub.5 placed in the 5.sup.th column of the filter bank G (described below) will not be utilized with any of the pixels of the degraded frame **602**, then all c.sub.5.sup.p,q output by the ML model **604** will be zero.

(67) At **706**, the filter bank G (i.e., the filter bank **610** of FIG. **6**) is obtained. In an example, the filter bank can be obtained by placing the G.sub.side and the G.sub.fixed matrices side-by-side, as shown in equation (5).

$$(68) \quad G = [G_side \ .\text{Math.} \ G_fixed] \quad (5)$$

(69) The encoder can determine the most beneficial side information regarding side filters (and side combining scalars, which are described below) using any number of techniques. For example, using peak signal-to-noise ratio (PSNR) and/or any error metric (e.g., sum of the mean squared error, sum of absolute differences error), the encoder can determine the side filters that are the best at minimizing the error metric. To illustrate, and without limitations, in a case that the error is greater than a threshold, the encoder may increase the side information, such as by sending additional filters instead of a current number of filters. The encoder may further modify the combining scalars to determine whether the modifications reduce the errors. In an example, the encoder may add one (or other values) to at least some of the combining scalars obtained from the ML model. The encoder may determine how much side information to transmit based on a rate-distortion calculation. For example, the encoder may obtain the optimal side information to transmit ignoring any rate limitations and then determine the subset of the side information to transmit based on an available number of bits balanced with the distortion reduction that results therefrom.

(70) Referring again to **608** of FIG. **6**, a filter f.sup.p,q for the pixel at location (p, q) can be obtained using equation (6). As such, from the N filters of the filter bank G, one filter f.sup.p,q is obtained.

$$(71) \quad f^{p,q} = Gc^{p,q} \quad (6)$$

(72) To illustrate, and without limitations, assume that there are 2 side filters, g.sub.1 and g.sub.2, one fixed filter g.sub.3, and that the window size K is 5. As such, f.sup.p,q for the pixel at (p=5, q=67) can be obtained using equation (7).

$$(73) \quad f^{5,67} = \begin{bmatrix} g_{1,1} & g_{2,1} & g_{3,1} \\ g_{1,2} & g_{2,2} & g_{3,2} \\ g_{1,3} & g_{2,3} & g_{3,3} \\ g_{1,4} & g_{2,4} & g_{3,4} \\ g_{1,5} & g_{2,5} & g_{3,5} \end{bmatrix} \begin{bmatrix} c_1^{5,67} \\ c_2^{5,67} \\ c_3^{5,67} \end{bmatrix} \quad (7)$$

(74) At **612**, each pixel value x(p, q) of the degraded frame **602** is filtered using its respective filter f.sup.p,q to obtain a corresponding restored pixel value {circumflex over (x)}(p, q) of a restored frame **614**. In an example, and in the case of K=k.sup.2 filters defined within a k×k window, the linear convolution operation of equation (8) can be used.

$$(75) \quad \hat{x}(p, q) = \sum_{m,n=0}^{k-1} x(p+m-\frac{k}{2}, q+n-\frac{k}{2}) f^{p,q}(m * k + n + 1) \quad (8)$$

(76) Returning briefly to the ML model **604**, the ML model **604** can be trained to minimize the error between restored frames obtained using equation (8) and the corresponding source (i.e., original) frames as the ML model (i.e., the combining scalars output therefrom) attempts to make the restored frame as close to the original frame as possible. The error can be backpropagated through the ML model **604** to adjust the weights of the ML model.

(77) In some implementations, the side filters can be transmitted at the frame level, as already alluded to. In another example, the side filters can be transmitted at the block level. As such, respective side filters may be transmitted for at least some of the blocks of the degraded frame. In an example, the side filters may be transmitted for a group-of-pictures (GOP). As such, the same side filters G.sub.side are used for each frame of the GOP.

(78) As already described with respect to equation (5), the side filters G.sub.side are used to expand the filter bank of the fixed filters G.sub.fixed. In another example, the side filters G.sub.side can be transmitted

as differential filters on top of the fixed filters, as shown in equation 5'.

$$(79) \quad G = G_{\text{fixed}} + G_{\text{side}} \quad (5')$$

(80) As such, each combining scalars c.sup.p,q output by the ML model **604** includes a number of scalars that is equal to the number of filters in the fixed filters G.sub.fixed. Alternatively, the ML model **604** may output more scalars than the number of fixed filters and any scalars not corresponding to the fixed filters can be ignored.

(81) One or more syntax elements of the bitstream can be used by the decoder to determine how the side filters G.sub.side are to be used. For example, a first value of the syntax element can indicate that the side filters are expanding filters (e.g., equation (5)), a second value of the syntax can indicate that the side filters are differential filters (e.g., equation (5')), and so on. The encoder may determine to use and transmit expanding filters in a case where the fixed filters do not produce a restored frame that is sufficiently close to the original frame and, as such, the encoder transmits more filters to be used by the decoder. The encoder may determine that additional slight improvements can be obtained in addition to using the fixed filters and, as such, the encoder transmits adjustments to some (e.g., a few) of the tap values of some of the fixed filters.

(82) In an example, the encoder may transmit, and the decoder may use, differentials $\delta_{\text{sup.p,q}}$ for the combining scalars. The encoder may determine that transmitting the differential combining scalars $\delta_{\text{sup.p,q}}$ can further improve the restored frame **614**. The ML model **604** may be trained on many video sequences (i.e., frames of the video sequences). However, the current video sequence may be sufficiently different from the training images and the encoder may determine that the output of the ML model **604** can be improved upon. As such, the encoder may determine to transmit, for at least some of the pixels of the degraded frame **602**, respective updates. For most of the pixels, the updates may be zero values (i.e., $\delta_{\text{i}\{\text{circumflex over ()}\}(p,q)=0}$). However, for other pixels, the respective updates will not be zero (i.e., $\delta_{\text{sub.i}\{\text{sup.p,q}\}\neq 0}$). The differential combining scalars $\delta_{\text{sup.p,q}}$ can be used to update the combining scalars c.sup.p,q output by the ML model **604** to obtain updated combining scalars $\{\text{tilde over (c)}\}_{\text{sup.p,q}}$, as shown in equation (9). The updated combining scalars $\{\text{tilde over (c)}\}_{\text{sup.p,q}}$ are then used to obtain the pixel-specific filters, as shown in equation (6').

$$(83) \quad \tilde{c}^{p,q} = c^{p,q} + \delta_{\text{sup.p,q}} \quad (9) \quad f^{p,q} = G\tilde{c}^{p,q} \quad (6')$$

(84) In another example, pixel-specific filters can be derived for groups of pixels rather than, as already described, for individual pixels. As such, each pixel in the group can be filtered with the group-specific filter. For example, a group of pixels can be a B×B block of pixels and one filter f is obtained and used for each pixel of the block. This results in one filter per B*B pixels therewith reducing derivation-related calculations.

(85) In another example, the ML model **604** can be trained to output for each pixel location, in addition to the N combining scalars c.sup.p,q, a pixel offset x.sub.offet.sup.p,q. The pixel offset x.sub.offet.sup.p,q can be added to the filtered pixel $\{\text{circumflex over (x)}\}(p,q)$. As such, the pixel values of the restored frame **614** can be given by equation (8').

$$(86) \quad \tilde{x}'(p,q) = \tilde{x}(p,q) + x_{\text{offet}}^{p,q} \quad (8')$$

(87) FIG. **8** is a flowchart of a technique **800** for restoring a degraded frame. The technique **800** can be implemented in a decoder such as the decoder **500** and can be implemented, for example, as a software program that can be executed by computing devices (e.g., apparatuses) such as receiving station **106**. The software program can include machine-readable instructions that can be stored in a memory (e.g., a non-transitory computer-readable storage medium) such as the memory **204** or the secondary storage **214**, and that can be executed by a processor, such as CPU **202**, to cause the computing device to perform the technique **800**. In at least some implementations, the technique **800** can be performed in whole or in part by the reconstruction stage **510** of the decoder **500** of FIG. **5**. In at least some implementations, the technique **800** can be performed in whole or in part by the loop filtering stage **512** of the decoder **500** of FIG. **5**.

(88) The technique **800** can be implemented using specialized hardware or firmware. Some computing devices can have multiple memories, multiple processors, or both. The steps or operations of the technique **800** can be distributed using different processors, memories, or both. Use of the terms “processor” or “memory” in the singular encompasses computing devices that have one processor or one memory as well as devices that have multiple processors or multiple memories that can be used in the performance of some or all of the recited steps.

(89) At **802**, a filter bank that includes filters is obtained. The filter bank can be as described with respect to the filter bank **G** of FIG. **6**. In an example, the filter bank can include side filters decoded from a compressed bitstream. The side filters can be as described with respect to **G.sub.side** and can be obtained as described with respect to **702** of FIG. **7**. In an example, the filter bank can further include fixed filters available at a decoder, as described with respect to **G.sub.fixed**, which can be obtained as described with respect to **704** of FIG. **7**. As such, the filter bank can be obtained as described with respect to equation (5).

(90) At **804**, respective sets of combining scalars for combining the filters of the filter bank can be obtained for pixels of a degraded frame. In an example, a respective set of combining filters can be obtained for each pixel of the degraded frame. Each set of combining scalars can be as described with respect to **c.sup.p,q** above. At **806**, respective pixel-specific filters can be obtained for the pixels of the degraded frame by combining the filters of the filter bank using the respective sets of combining scalars, as described with respect to equation (6). At **808**, a restored frame is obtained by filtering the pixels of the degraded frame using the respective pixel-specific filters. The restored frame can be the restored frame **614** of FIG. **6**, which can be obtained using equation (8).

(91) Returning again to FIG. **6**, in an example, look-up-tables (LUTs) can be used to significantly accelerate computations in an area **616**. As can be appreciated, the process described with respect to FIG. **6** performs many calculations including those performed by the ML model **604**. To reduce the computational complexity, LUTs can be used to look up, rather than perform to obtain, calculation results (or approximations thereof). As can also be appreciated, the LUTs cannot be infinitely large to account for all possible input values. As such, the operands of operations (e.g., multiplications, convolutions) performed within the area **616** may be quantized to nearest values for lookup in the LUTs.

(92) FIG. **9** is a flowchart of an example of a technique **900** for filtering with pixel-specific filters obtained using a LUT. The technique **900** can be implemented by a loop filtering stage, such as the loop filtering stage **416** of FIG. **4** or the loop filtering stage **512** of FIG. **5**. Filtering using combining scalars can be implemented by a reconstruction stage, such as the reconstruction stage **510** of FIG. **5** or the reconstruction stage **414** of FIG. **4**.

(93) The technique **900** obtains, for a degraded frame **902**, pixel-specific filters **f.sup.p,q** that are used, at **912** (which can be or be similar to **612** of FIG. **6**), to filter the pixels of a degraded frame **902** (which can be or be similar to the degraded frame **602** of FIG. **6**) to obtain a restored frame **914** (which can be or be similar to the restored frame **614** of FIG. **6**).

(94) An ML model **904** receives a degraded frame **902**. Whereas the ML model **604** of FIG. **6** outputs, for pixels of the degraded frame **602**, respective combining scalars **c.sup.p,q** that are vectors of values, the ML model **904** outputs an index **t.sup.p,q** per pixel. The index **t.sup.p,q** is used, at **906**, as an index into a LUT **908** to obtain pre-computed (and stored) combining scalars **c.sub.fixed.sup.p,q**. **C.sub.fixed** (i.e., “fixed combining scalars”) refers to the combining scalars obtained from the LUT **908** using all of the **t.sup.p,q** indexes. The ML model **904** may also output a filter modifier **λ.sup.p,q** for each pixel, which is explained further below.

(95) As described above, the compressed bitstream **420** may include side filters **G.sub.side**. In such a case, at **916**, the technique **900** decodes the side filters **G.sub.side** from the compressed bitstream **420**. At **918**, the filter bank **G** can be obtained by combining the side filters **G.sub.side** (e.g., if any) and the fixed filters **G.sub.fixed**, as described above, such as with respect to equation (5) or equation (5’).

(96) As also described above, the compressed bitstream **420** may include information that may be used to update the combining scalars. As such, at **920**, the combining scalar side information **C.sub.side** may be decoded from the compressed bitstream. At **922**, the fixed combining scalars **C.sub.fixed** can be combined with the scalar side information **C.sub.side** (if any) to obtain a combining scalar matrix **C**. **C.sub.side** and **C.sub.fixed** can be combined in an expanding, a differential way, or some other way, depending on a syntax value of the compressed bitstream. The combining scalar matrix **C** is used with the filter bank **G** (i.e., $F=GC$) along with the indexes **t.sup.p,q** as input to the LUT **908** to obtain pixel-specific stored filters **F(:, t.sup.p,q)**.

(97) At **910**, the pixel-specific stored filters **F(:, t.sup.p,q)** can be modified using the filter modifier **λ.sup.p,q** to obtain the pixel-specific filter **f.sup.p,q**, which is used to filter the pixel at location (p, q) of the degraded frame **902**.

(98) To restate, whereas the ML model **604** performs complicated computations to generate **N** numbers per pixel (i.e., **c.sup.p,q**), the ML model **904** generates one number **t.sup.p,q** per pixel, which can be

immediately used to obtain from the LUT **908** a pixel-specific stored filters $F(:, t.\text{sup.p}, q)$. Depending on the structure of the LUT, the pixel-specific stored filters $F(:, t.\text{sup.p}, q)$ may be stored as rows or columns in the LUT **908**. The filter modifier $\lambda.\text{sup.p}, q$ is then used to modify the stored pixel-specific filter.

(99) As mentioned above, in some examples, the combining scalars $c.\text{sup.p}, q$ can be obtained using techniques other than a neural network.

(100) Let $p.\text{sub.k}, l(L \times 1)$ denote a vector formed by a patch around a pixel at (k, l) of a restoration unit. Additionally, assume that H (of size $L \times L$) is an orthonormal transform and $T > 0$ is a given threshold. Well-known orthonormal-transform-and-hard-threshold-based denoising of this patch reconstructs the vector $\{\text{circumflex over (p)}\}.\text{sub.k}, l$ as given by equation (10), where $h.\text{sub.i}, i=1, \dots, L$ are the columns of H .

(101)

$$\hat{p}_{k,l} = p_{k,l} - \text{Math.}_{\{i: \text{Math. } h_i^T p_{k,l} \text{ Math. } < T\}} h_i^T p_{k,l} h_i = p_{k,l} - (\text{Math.}_{\{i: \text{Math. } h_i^T p_{k,l} \text{ Math. } < T\}} h_i h_i^T) p_{k,l} \quad (10)$$

(102) Considering one of the components (i.e., n) of $p.\text{sub.k}, l$ that corresponds to the pixel (k, l) within the patch, the reconstructed pixel $\{\text{circumflex over (p)}\}.\text{sub.k}, l(n)$ can be calculated using equation (11).

(103) $\hat{p}_{k,l}(n) = p_{k,l}(n) - (\text{Math.}_{\{i: \text{Math. } h_i^T p_{k,l} \text{ Math. } < T\}} h_i(n) h_i^T) p_{k,l} = p_{k,l}(n) - f_n^T p_{k,l} \quad (11)$

where $f_n = \text{Math.}_{\{i: \text{Math. } h_i^T p_{k,l} \text{ Math. } < T\}} h_i(n) h_i^T$

(104) In equation (11), $f.\text{sub.n}$ can be considered a pixel-adaptive filter. The pixel-adaptive filter $f.\text{sub.n}$ can be realized by steps including evaluating features $|h.\text{sub.i}.\text{sup.Tp}.\text{sub.k}, l|$, comparing these features to thresholds, forming the appropriate filter using equation (11), and finally forming the reconstructed pixel value.

(105) Considering equation (11), the pixel-adaptive filter can be observed to be constructed (e.g., put together) using L features, where each feature independently contributes to an incremental component of the filter. Considering, by way of a non-limiting example, a patch size of 7×7 pixels, $L=49$ may lead to substantial complexity at an encoder and a decoder.

(106) Implementations according to this disclosure can use a few number of features F . In an example, $F=4$. The features F are such that they are used jointly in determining the pixel-adaptive filter for a pixel. In an example, the features can be quantized and combined for use in a lookup table (LUT) of filters. The features F can be considered to be equivalent to the combining scalars $c.\text{sup.p}, q$ described above; and the process described with respect to obtaining the features can be considered to be a simplification of the neural network using to obtain the combining scalars $c.\text{sup.p}, q$.

(107) Specifically, quantized features $w.\text{sub.i}$ using equation (12) can be obtained. The quantized features $w.\text{sub.i}$ can be used to determine the pixel-adaptive filter using equation (13).

(108) $w_i = (\text{Math. } g_i^T p_{k,l} \text{ Math. } - T_i), i = 1, \text{Math. }, 4 \quad (12) \quad f_{k,l} = LUT(w_1, w_2, w_3, w_4) \quad (13)$

(109) In equation (12),  is a quantization function, $g.\text{sub.i}$ are feature generation projections, $T.\text{sub.i}$ are thresholds, and $|g.\text{sub.i}.\text{sup.Tp}.\text{sub.k}, l|$ can be considered to be the features F corresponding to the combining scalars. In equation (13), $f.\text{sub.k}, l$ is the pixel adaptive filter to be applied to the pixel (k, l) . The thresholds $T.\text{sub.i}$ can be considered to be similar to regularization parameter that may typically be used when training a neural network. The filter $f.\text{sub.k}, l$ is obtained from the lookup table using the function LUT, which takes the quantized features $w.\text{sub.i}$ as input.

(110) The quantization function  used in equation (12) can depend on the number of features, the number of entries in the LUT, or both. To illustrate, assume that the LUT includes 4096 entries (i.e., filters). As four features are used in equation (12), the quantization function  can be an eight-level quantizer. As such, each of the feature values can be quantized to a value between 0 and 7, therewith obtaining a total of $8 \times 8 \times 8 \times 8 = 8.\text{sup.4} = 4096$ possible combination of values. Obtaining a value between 0 and 4095 for a pixel can essentially be understood to, in effect, classify the pixel or, more accurately, the neighborhood of the pixel into one of the 4096 possibilities.

(111) In another example, the LUT table may include 256 entries and LUT(.) function may map an input value that is in the range $[0, 4095]$ to one of the 256 filters. The quantized feature values obtained using equation (12) can be used in equation (13) to obtain a filter based on the feature values. The quantized feature values may be combined in any number of ways. In an example, each of feature values may be quantized into 3 bits therewith obtaining a 12-bit value that can be obtained by concatenating the bits of the quantized feature values. The 12-bit value can be used as an input to the LUT(Math.) of equation (13).

(112) The filters of the LUT can be obtained using offline training. A training set of patches can be used to optimize the filters over a range of quality levels. In an example, the training set can be formed using 30 frames each from primarily 720p and 1080p video sequences. The thresholds $T_{sub,i}$ may be optimized for each quality level. In an example, and in order to further streamline computations, simple projections $g_{sub,i}$ can be formed by the one-dimensional $[-1, 2, -1]$ gradient filter configured on horizontal, vertical, diagonal, and anti-diagonal directions and computing respective averages over each patch. More generally, projections (e.g., filters) that can highlight high frequency areas (e.g., pixel neighborhoods) can be used. For example, Laplacian-like operators, which can act as second derivatives, can be used.

(113) In another example, simple projections can be used (e.g., combined) to obtain more complex features. Gradient filters using a small number of taps can be operated (e.g., applied), with respect to a pixel, in at least some of the horizontal, vertical, diagonal, and anti-diagonal directions. In an example, the number of taps can be 3. In an example, the filter weights of all of the filters can be $[-1, 2, -1]$. However, that need not be the case and different directions can use different weights, different number of taps, or both.

(114) For at least some (e.g., each) of the pixels, an averaged magnitude of each filter leads to a classification feature. In some examples, thresholds can be subtracted from the classification features and the results can be used to consult a filter lookup-table (LUT). The thresholds can be imperially derived. The LUT in turn yields an origin-symmetric non-separable filter. The derived filter is then used to obtain the filtered output at that pixel. The encoder and decoder perform the same set of calculations in deriving the filter. The decoder only performs the calculation on RUs where the mode is signaled.

(115) FIG. 10 illustrates an example 1000 of obtaining magnitude feature values. FIG. 10 includes a restoration unit 1001. The restoration unit 1001 is shown as being of size 11×11 pixels. However, the disclosure is not so limited and the restoration unit 1001 can be smaller or larger. In an example, the restoration unit 1001 can have a size of 256×256 pixels for a luma component and a size of 128×128 pixels for chroma component. FIG. 10 illustrates that filters are applied at pixels (i.e., at each of the pixels) of the restoration unit 1001. One such pixel is a pixel 1002, which may be at a location (k, l) within the restoration unit 1001. FIG. 10 illustrates that 4 filters are applied to (or operated at) the pixel 1002 to obtain magnitude feature values. The filters are or include a horizontal filter 1004, a vertical filter 1006, a diagonal filter 1008, and an anti-diagonal filter 1010. While FIG. 10 is described with respect to a number of filters N that is equal to 4, the disclosure is not so limited. N can be larger or smaller than 4. Each of the filters is shown as including the same number of taps (i.e., 3) and the same weights (i.e., $[-1, 2, -1]$). However, that need not be the case. More taps per filter can be used, different weights per filter can be used, or a combination thereof. Each of the filters is operated over a patch 1012 (i.e., a window) around (e.g., surrounding) the pixel 1002 to obtain, respectively, magnitude feature values 1014-1020; namely $f_{sub,n=0}(k, l)$, $f_{sub,n=1}(k, l)$, $f_{sub,n=2}(k, l)$, and $f_{sub,n=3}(k, l)$. Again, four (i.e., $n=0, \dots, 3$) filters are illustratively used; but more or fewer filters can be used. Thus, in the general case, $n=0, \dots, N-1$.

(116) Table I illustrates a pseudocode for obtaining the magnitude feature values 1014-1020. Other algorithms for obtaining the magnitude feature values 1014-1020 are possible. The magnitude feature values 1014-1020 are calculated at lines 6-9, respectively, of Table I. In Table I, ru is a 2-dimensional array that holds pixel values of a restoration unit, such as the restoration unit 1001 of FIG. 10. The pseudocode of Table I calculates the magnitude feature values 1014-1020 for the pixel at location (k, l) of the restoration unit.

(117) TABLE-US-00001
 TABLE I
 1 $base_value = 2 * ru(k, l)$
 2 $horizontal_diff = ru(k - 1, l) + ru(k + 1, l)$
 3 $vertical_diff = ru(k, l - 1) + ru(k, l + 1)$
 4 $anti_diagonal_diff = ru(k - 1, l + 1) + ru(k + 1, l - 1)$
 5 $diagonal_diff = ru(k - 1, l - 1) + ru(k + 1, l + 1)$
 6 $f_0 = abs(base_value - horizontal_diff)$
 7 $f_1 = abs(base_value - vertical_diff)$
 8 $f_2 = abs(base_value - diagonal_diff)$
 9 $f_3 = abs(base_value - anti_diagonal_diff)$

(118) After obtaining the magnitude feature values 1014-1020 for each pixel of the restoration unit 1001, classification features $af_{sub,n}(k, l)$, $n=0, \dots, 3$, are obtained by averaging the corresponding magnitude feature values $f_{sub,n}(k, l)$, $n=0, \dots, 3$ in a window 1022 (i.e., a neighborhood) centered at the pixel (k, l) . While the window 1022 is shown as being a 5×5 neighborhood, that need not be the case and smaller or larger neighborhood $P \times Q$ can be used. In an example, the neighborhood is a square and, as such, $P=Q$. The classification features can be obtained using equation (14).

$$(119) \quad af_n(k, l) = \frac{\text{Math}_{p=0}^{P-1} \cdot \text{Math}_{q=0}^{Q-1} f_n(i, j)}{P * Q} \quad (14)$$

(120) To restate, each of the magnitude feature values are averaged, as a heuristic, over a window (in this example, a 5×5 window). The classification features $af_{sub.n}(k, l)$ can be thought as being equivalent to the quantities $|g_{sub.i.sup.Tp.sub.k,l}|$ of equation (12). Each of the $af_{sub.n}(k, l)$, $n=0, \dots, 3$ can be quantized to obtain quantized features $cf_{sub.n}(k, l)$, $n=0, \dots, 3$ using equation (15). The quantized features $c_{sub.fn}(k, l)$ of equation (15) can be thought of as being equivalent to the quantized features $w_{sub.i}$ obtained in equation (12).

$$(121) \quad cf_n(k, l) = \text{quantize}(af_n(k, l) - \text{threshold}_n) \quad (15)$$

(122) The $cf_{sub.n}(k, l)$ can be combined, such as described with respect to equation (13), and the output be used as inputs into a lookup table to obtain a non-separable filter $f_{sub.k,l}$ that can be used to filter the pixel at (k, l) .

(123) FIG. 11 is an illustration of a portion of a lookup table 1100. The lookup table includes pixel-adaptive filters, illustrated as squares in the FIG. 11. Filters 1102, 1104, and 1106 may correspond to inputs, or be returned by a lookup function that takes as inputs, the values 0, 1, and 2, respectively, corresponding to feature values as described above. At least some of the same filters may be used at different quality levels. As such, it can be observed that similar filters with light variations that address different levels of quantization noise are included in the lookup table 1100. Additionally, as can be observed, some of the filters have directional structure and frequency content (as illustrated by filters 1108-1110, amongst others) that would at best be difficult, if not impossible, to realize using separable versions of such filters.

(124) FIG. 12 is a flowchart of a technique 1200 for restoring a degraded frame. The technique 1200 can be implemented in a decoder such as the decoder 500 and can be implemented, for example, as a software program that can be executed by computing devices (e.g., apparatuses) such as receiving station 106. The software program can include machine-readable instructions that can be stored in a memory (e.g., a non-transitory computer-readable storage medium) such as the memory 204 or the secondary storage 214, and that can be executed by a processor, such as CPU 202, to cause the computing device to perform the technique 1200. In at least some implementations, the technique 1200 can be performed in whole or in part by the reconstruction stage 510 of the decoder 500 of FIG. 5. In at least some implementations, the technique 1200 can be performed in whole or in part by the loop filtering stage 512 of the decoder 500 of FIG. 5. The technique 1200 can be performed for pixels (e.g., each of the pixels) of a restoration unit. FIG. 12 is described with respect to performing the technique 1200 for a pixel, which may be the pixel 1002 of FIG. 10 of a degraded frame.

(125) At 1202, magnitude features (i.e., magnitude feature values) are obtained based on a window (a first window) centered at the pixel. A cardinality N of the magnitude features is at least 1. In an example, the magnitude features can be the features F (i.e., $|g_{sub.i.sup.Tp.sub.k,l}|$) described with respect to equation (12). The window can be the patch 1012 of FIG. 10. In an example, the cardinality of the magnitude features is 4. That is, four features are obtained for the pixel. In an example, the magnitude features can be magnitude feature values 1014-1020 of FIG. 10. As such, the magnitude feature may be obtained as described with respect to Table I.

(126) In an example, the magnitude features can be obtained using filters that include at least two of a horizontal filter, a vertical filter, a diagonal filter, or an anti-diagonal filter, which may be as described with respect to FIG. 10. As such, in an example, each of the filters can be a 3-tap filter that uses the weights $[-1, 2, 1]$.

(127) At 1204, the magnitude features are used to obtain a pixel-adaptive filter. In an example, the magnitude features are used to obtain a pixel-adaptive filter as described with respect to equations (12)-(13). In another example, the magnitude features are used to obtain a pixel-adaptive filter as described with respect to FIG. 10 and equations (14)-(15). As such, using the magnitude features to obtain the pixel-adaptive pixel can include obtaining, for the pixel, N classification features. Each of the N classification features corresponds an average of respective magnitude features of pixels of a second window that is centered at the pixel. The second window can be the window 1022 of FIG. 10. In an example, the first window can have a size 3×3 and the second window can have a size of 5×5.

(128) In an example, using the magnitude features to obtain the pixel-adaptive pixel further can include using the magnitude features to obtain the pixel-adaptive filter from a lookup table. In an example, using the magnitude features to obtain the pixel-adaptive filter from a lookup table can include quantizing at least some of the N classification features and using the at least some of the N classification features to obtain the pixel-adaptive filter from the lookup table.

(129) At **1206**, the pixel-adaptive filter can be applied to the pixel to obtain a pixel of the restored frame, as described above with respect to FIGS. **6-9**. In an example, obtaining the pixel-adaptive filter, and as described above, can include, combining the pixel-adaptive filter obtained from the lookup table with side filters obtained using side information received from an encoder. The resulting pixel-adaptive filter can be applied to the pixel to obtain the restored pixel of the restored frame.

(130) Returning briefly to FIGS. **6** and **9**, the ML model **604** and/or the ML model **904** can each be any type of ML model that is capable of being trained to receive a video frame and output a vector of scalar values or other filter-related information, as described herein. In an example, the ML model **604** and the ML model **904** can each be a neural network. In an example, the neural network can be a deep-learning convolutional ML model (CNN). In a CNN, a feature extraction portion typically includes a set of convolutional operations, which is typically a series of filters that are used to filter an input (e.g., an image) based on a filter (typically a square of size l , without loss of generality). For example, in machine vision (i.e., the processing of an image of a patient's room), these filters can be used to find features in an input image. The features can include, for example, edges, corners, endpoints, and so on. As the number of stacked convolutional operations increases, later convolutional operations can find higher-level features.

(131) In the CNN, a classification portion is typically a set of fully connected layers. The fully connected layers can be thought of as looking at all the input features of an image in order to generate a high-level classifier. Several stages (e.g., a series) of high-level classifiers eventually generate the desired classification output.

(132) As mentioned, a typical CNN network is composed of a number of convolutional operations (e.g., the feature-extraction portion) followed by a number of fully connected layers. The number of operations of each type and their respective sizes is typically determined during a training phase of the machine learning. As a person skilled in the art recognizes, additional layers and/or operations can be included in each portion. For example, combinations of Pooling, MaxPooling, Dropout, Activation, Normalization, BatchNormalization, and other operations can be grouped with convolution operations (i.e., in the features-extraction portion) and/or the fully connected operation (i.e., in the classification portion). The fully connected layers may be referred to as Dense operations. As a person skilled in the art recognizes, a convolution operation can use a SeparableConvolution2D or Convolution2D operation.

(133) A convolution layer can be a group of operations starting with a Convolution2D or SeparableConvolution2D operation followed by zero or more operations (e.g., Pooling, Dropout, Activation, Normalization, BatchNormalization, other operations, or a combination thereof), until another convolutional layer, a Dense operation, or the output of the CNN is reached. A convolution layer can use (e.g., create, construct, etc.) a convolution filter that is convolved with the layer input to produce an output (e.g., a tensor of outputs). A Dropout layer can be used to prevent overfitting by randomly setting a fraction of the input units to zero at each update during a training phase. A Dense layer can be a group of operations or layers starting with a Dense operation (i.e., a fully connected layer) followed by zero or more operations (e.g., Pooling, Dropout, Activation, Normalization, BatchNormalization, other operations, or a combination thereof) until another convolution layer, another Dense layer, or the output of the network is reached. The boundary between feature extraction based on convolutional networks and a feature classification using Dense operations can be marked by a Flatten operation, which flattens the multidimensional matrix from the feature extraction into a vector.

(134) In a typical CNN, each of the convolution layers may consist of a set of filters. While a filter is applied to a subset of the input data at a time, the filter is applied across the full input, such as by sweeping over the input. The operations performed by this layer are typically linear/matrix multiplications. The activation function may be a linear function or non-linear function (e.g., a sigmoid function, an arcTan function, a tanH function, a ReLu function, or the like).

(135) Each of the fully connected operations is a linear operation in which every input is connected to every output by a weight. As such, a fully connected layer with N number of inputs and M outputs can have a total of $N \times M$ weights. As mentioned above, a Dense operation may be generally followed by a non-linear activation function to generate an output of that layer.

(136) The aspects of encoding and decoding described above illustrate some encoding and decoding techniques. However, it is to be understood that encoding and decoding, as those terms are used in the claims, could mean compression, decompression, transformation, or any other processing or change of data.

(137) The words "example" or "implementation" are used herein to mean serving as an example, instance,

or illustration. Any aspect or design described herein as “example” or “implementation” is not necessarily to be construed as preferred or advantageous over other aspects or designs. Rather, use of the words “example” or “implementation” is intended to present concepts in a concrete fashion. As used in this application, the term “or” is intended to mean an inclusive “or” rather than an exclusive “or”. That is, unless specified otherwise, or clear from context, “X includes A or B” is intended to mean any of the natural inclusive permutations. That is, if X includes A; X includes B; or X includes both A and B, then “X includes A or B” is satisfied under any of the foregoing instances. In addition, the articles “a” and “an” as used in this application and the appended claims should generally be construed to mean “one or more” unless specified otherwise or clear from context to be directed to a singular form. Moreover, use of the term “an implementation” or “one implementation” throughout is not intended to mean the same embodiment or implementation unless described as such.

(138) Implementations of transmitting station **102** and/or receiving station **106** (and the algorithms, methods, instructions, etc., stored thereon and/or executed thereby, including by encoder **400** and decoder **500**) can be realized in hardware, software, or any combination thereof. The hardware can include, for example, computers, intellectual property (IP) cores, application-specific integrated circuits (ASICs), programmable logic arrays, optical processors, programmable logic controllers, microcode, microcontrollers, servers, microprocessors, digital signal processors or any other suitable circuit. In the claims, the term “processor” should be understood as encompassing any of the foregoing hardware, either singly or in combination. The terms “signal” and “data” are used interchangeably. Further, portions of transmitting station **102** and receiving station **106** do not necessarily have to be implemented in the same manner.

(139) Further, in one aspect, for example, transmitting station **102** or receiving station **106** can be implemented using a general-purpose computer or general-purpose processor with a computer program that, when executed, carries out any of the respective methods, algorithms and/or instructions described herein. In addition, or alternatively, for example, a special purpose computer/processor can be utilized which can contain other hardware for carrying out any of the methods, algorithms, or instructions described herein.

(140) Transmitting station **102** and receiving station **106** can, for example, be implemented on computers in a video conferencing system. Alternatively, transmitting station **102** can be implemented on a server and receiving station **106** can be implemented on a device separate from the server, such as a hand-held communications device. In this instance, transmitting station **102** can encode content using an encoder **400** into an encoded video signal and transmit the encoded video signal to the communications device. In turn, the communications device can then decode the encoded video signal using a decoder **500**. Alternatively, the communications device can decode content stored locally on the communications device, for example, content that was not transmitted by transmitting station **102**. Other transmitting station **102** and receiving station **106** implementation schemes are available. For example, receiving station **106** can be a generally stationary personal computer rather than a portable communications device and/or a device including an encoder **400** can also include a decoder **500**.

(141) Further, all or a portion of implementations of the present disclosure can take the form of a computer program product accessible from, for example, a tangible computer-usable or computer-readable medium. A computer-usable or computer-readable medium can be any device that can, for example, tangibly contain, store, communicate, or transport the program for use by or in connection with any processor. The medium can be, for example, an electronic, magnetic, optical, electromagnetic, or a semiconductor device. Other suitable mediums are also available.

(142) The above-described embodiments, implementations and aspects have been described to allow easy understanding of the present disclosure and do not limit the present disclosure. On the contrary, the disclosure is intended to cover various modifications and equivalent arrangements included within the scope of the appended claims, which scope is to be accorded the broadest interpretation to encompass all such modifications and equivalent structure as is permitted under the law.

Claims

1. A method for obtaining a restored frame from a degraded frame, comprising: obtaining, for a pixel of the degraded frame, magnitude features based on a first window centered at the pixel, wherein a cardinality N of the magnitude features is at least 1; using the magnitude features to obtain a pixel-adaptive filter, wherein

using the magnitude features to obtain the pixel-adaptive pixel comprises: obtaining, for the pixel, N classification features, wherein each of the N classification features corresponds to an average of respective magnitude features of pixels of a second window that is centered at the pixel; and applying the pixel-adaptive filter to the pixel to obtain a pixel of the restored frame.

2. The method of claim 1, wherein a cardinality of the magnitude features is 4.

3. The method of claim 1, wherein the magnitude features are obtained using filters comprising at least two of a horizontal filter, a vertical filter, a diagonal filter, or an anti-diagonal filter.

4. The method of claim 1, wherein each of the filters is a 3-tap filter that uses weights $[-1, 2, 1]$.

5. The method of claim 1, wherein the second window is different from the first window.

6. The method of claim 1, wherein the first window has a size of 3×3 , and the second window has a size of 5×5 .

7. The method of claim 1, wherein using the magnitude features to obtain the pixel-adaptive pixel further comprises: using the magnitude features to obtain the pixel-adaptive filter from a lookup table.

8. The method of claim 1, wherein using the magnitude features to obtain the pixel-adaptive filter comprises: quantizing at least some of the N classification features; and using the at least some of the N classification features to obtain the pixel-adaptive filter from a lookup table.

9. A device, comprising: a processor, the processor configured to: obtain, for a pixel of a degraded frame, magnitude features based on a first window centered at the pixel, wherein a cardinality N of the magnitude features is at least 1; use the magnitude features to obtain a pixel-adaptive filter, wherein to use the magnitude features to obtain the pixel-adaptive pixel comprises to: obtain, for the pixel, N classification features, wherein each of the N classification features corresponds to an average of respective magnitude features of pixels of a second window that is centered at the pixel; and apply the pixel-adaptive filter to the pixel to obtain a pixel of a restored frame.

10. The device of claim 9, wherein a cardinality of the magnitude features is 4.

11. The device of claim 9, wherein the magnitude features are obtained using filters comprising at least two of a horizontal filter, a vertical filter, a diagonal filter, or an anti-diagonal filter.

12. The device of claim 9, wherein each of the filters is a 3-tap filter that uses weights $[-1, 2, 1]$.

13. The device of claim 9, wherein the first window has a size of 3×3 , and the second window has a size of 5×5 .

14. The device of claim 9, wherein to use the magnitude features to obtain the pixel-adaptive pixel further comprises to: use the magnitude features to obtain the pixel-adaptive filter from a lookup table.

15. The device of claim 14, wherein to use the magnitude features to obtain the pixel-adaptive filter from a lookup table comprises to: quantize at least some of the N classification features; and use the at least some of the N classification features to obtain the pixel-adaptive filter from the lookup table.

16. A non-transitory computer-readable storage medium, comprising executable instructions that, when executed by a processor, cause performance of operations for obtaining a restored frame from a degraded frame, comprising: obtaining, for a pixel of the degraded frame, magnitude features based on a first window centered at the pixel, wherein a cardinality N of the magnitude features is at least 1: using the magnitude features to obtain a pixel-adaptive filter, wherein using the magnitude features to obtain the pixel-adaptive pixel comprises: obtaining, for the pixel, N classification features, wherein each of the N classification features corresponds to an average of respective magnitude features of pixels of a second window that is centered at the pixel; and applying the pixel-adaptive filter to the pixel to obtain a pixel of the restored frame.

17. The non-transitory computer-readable storage medium of claim 16, wherein the magnitude features are obtained using filters comprising at least two of a horizontal filter, a vertical filter, a diagonal filter, or an anti-diagonal filter.

18. The non-transitory computer-readable storage medium of claim 16, wherein the first window has a size of 3×3 , and the second window has a size of 5×5 .

19. The non-transitory computer-readable storage medium of claim 16, wherein using the magnitude features to obtain the pixel-adaptive pixel further comprises: using the magnitude features to obtain the pixel-adaptive filter from a lookup table.

20. The non-transitory computer-readable storage medium of claim 1, wherein using the magnitude features to obtain the pixel-adaptive filter comprises: quantizing at least some of the N classification features; and using the at least some of the N classification features to obtain the pixel-adaptive filter from a lookup table.

