

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12416965
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Chidambaram; Madhusudan et al.

Distribution of frequency budget in a processor

Abstract

In an embodiment, a processor may include a core domain comprising a plurality of processing cores, an uncore domain comprising an internal network, and a processor power management circuit. The processor power management circuit may be to: receive scalability hint values from the processing cores; determine a total core gain and a total uncore gain based at least in part on the scalability hint values; and distribute a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. Other embodiments are described and claimed.

Inventors: Chidambaram; Madhusudan (Bangalore, IN), Rudramuni; Vishwesh (Karnataka, IN), Zobel; Shmuel (Hinanit, IL), Gendler; Alexander (Kiriath Motzkin, IL)

Applicant: Intel Corporation (Santa Clara, CA)

Family ID: 1000008820442

Assignee: Intel Corporation (Santa Clara, CA)

Appl. No.: 17/644677

Filed: December 16, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20230195199 A1	Jun. 22, 2023

Publication Classification

Int. Cl.: G06F1/324 (20190101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (1/324)

References Cited**U.S. PATENT DOCUMENTS**

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
6205533	12/2000	Margolus	N/A	N/A
10678724	12/2019	ChoFleming et al.	N/A	N/A
10795853	12/2019	Nassif et al.	N/A	N/A
10891240	12/2020	Mathew et al.	N/A	N/A
11086816	12/2020	Fleming et al.	N/A	N/A
11200186	12/2020	Fleming et al.	N/A	N/A
2013/0111236	12/2012	Ananthakrishnan	713/320	G06F 1/3234
2014/0089688	12/2013	Man	713/300	G06F 1/3206
2014/0181538	12/2013	Shrall	713/300	G06F 3/14
2015/0277538	12/2014	Yasin et al.	N/A	N/A
2017/0031412	12/2016	Gendler et al.	N/A	N/A
2022/0026972	12/2021	Shah	N/A	G06F 1/324
2022/0100247	12/2021	Garg	N/A	G06F 1/324

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
2796961	12/2013	EP	N/A
3901767	12/2020	EP	N/A

OTHER PUBLICATIONS

European Patent Office, Extended European Search Report for Appl. No. 22203049.6 dated May 4, 2023 (10 pages). cited by applicant

Gupta et al., U.S. Appl. No. 17/448,600 entitled Gating of a Mesh Clock Signal in a Processor filed Sep. 23, 2021 (145 pages). cited by applicant

Primary Examiner: Suryawanshi; Suresh

Attorney, Agent or Firm: Nicholson De Vos Webster & Elliott LLP

Background/Summary**FIELD OF INVENTION**

(1) Embodiments relate generally to computer systems. More particularly, embodiments are related to controlling clock signals in computer systems.

BACKGROUND

(2) Advances in semiconductor processing and logic design have permitted an increase in the amount of logic that may be present on integrated circuit devices. As a result, computer system

configurations have evolved from a single or multiple integrated circuits in a system to multiple hardware threads, multiple cores, multiple devices, and/or complete systems on individual integrated circuits. Further, as the density of integrated circuits has grown, the power requirements for computing systems have also grown. As a result, there is a vital need for energy efficiency and conservation associated with integrated circuits.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram of a portion of a system in accordance with an embodiment of the present invention.
- (2) FIG. 2 is a block diagram of a processor in accordance with an embodiment of the present invention.
- (3) FIG. 3 is a block diagram of a multi-domain processor in accordance with another embodiment of the present invention.
- (4) FIG. 4 is an embodiment of a processor including multiple cores.
- (5) FIG. 5 is a block diagram of a micro-architecture of a processor core in accordance with one embodiment of the present invention.
- (6) FIG. 6 is a block diagram of a micro-architecture of a processor core in accordance with another embodiment.
- (7) FIG. 7 is a block diagram of a micro-architecture of a processor core in accordance with yet another embodiment.
- (8) FIG. 8 is a block diagram of a micro-architecture of a processor core in accordance with a still further embodiment.
- (9) FIG. 9 is a block diagram of a processor in accordance with another embodiment of the present invention.
- (10) FIG. 10 is a block diagram of a representative SoC in accordance with an embodiment of the present invention.
- (11) FIG. 11 is a block diagram of another example SoC in accordance with an embodiment of the present invention.
- (12) FIG. 12 is a block diagram of an example system with which embodiments can be used.
- (13) FIG. 13 is a block diagram of another example system with which embodiments may be used.
- (14) FIG. 14 is a block diagram of a representative computer system.
- (15) FIGS. 15A-15B are block diagrams of systems in accordance with embodiments of the present invention.
- (16) FIG. 16 is a block diagram illustrating an IP core development system used to manufacture an integrated circuit to perform operations according to an embodiment.
- (17) FIGS. 17A-17B are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to embodiments of the invention.
- (18) FIGS. 18A-D are block diagrams illustrating an exemplary specific vector friendly instruction format according to embodiments of the invention.
- (19) FIG. 19 is a block diagram of a register architecture according to one embodiment of the invention.
- (20) FIG. 20A is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the invention.
- (21) FIG. 20B is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the invention.

- (22) FIGS. 21A-B illustrate a block diagram of a more specific exemplary in-order core architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip.
- (23) FIG. 22 is a block diagram of a processor that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the invention.
- (24) FIGS. 23-24 are block diagrams of exemplary computer architectures.
- (25) FIG. 25 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the invention.
- (26) FIG. 26 is a block diagram of an example system in accordance with one or more embodiments.
- (27) FIG. 27 is a block diagram of an example system in accordance with one or more embodiments.
- (28) FIG. 28 is a flow diagram of an example method in accordance with one or more embodiments.
- (29) FIG. 29 is an illustration of an example operation in accordance with one or more embodiments.
- (30) FIG. 30 is an illustration of example formulas in accordance with one or more embodiments.
- (31) FIG. 31 is an illustration of an example table in accordance with one or more embodiments.
- (32) FIG. 32 is an illustration of an example table in accordance with one or more embodiments.
- (33) FIG. 33 is an illustration of an example table in accordance with one or more embodiments.
- (34) FIG. 34 is a flow diagram of an example method in accordance with one or more embodiments.
- (35) FIG. 35 is an illustration of an example storage medium in accordance with one or more embodiments.

DETAILED DESCRIPTION

- (36) Some computing devices may include internal networks to provide communication between various components. For example, some processors may include an internal network to allow included processing engines (e.g., “cores”) to transfer data with each other and/or other components of the processor (e.g., caches, memory, input/output devices, and so forth). In some examples, the internal network may be implemented as a mesh network or fabric that includes multiple interconnected nodes (referred to herein as “mesh stops”). The mesh stops may operate according to a shared clock signal (also referred to herein as a “clock frequency,” or as a “frequency”) that may be increased or decreased to vary the performance of the mesh stops. For example, the frequency of the clock signal may be increased when the mesh network is under a heavy workload, and may be decreased when the mesh network is under a light workload. However, increasing the clock signal will increase the power consumption of the mesh network, and therefore may result in inefficient power use when the increased performance is not needed.
- (37) In some processors, various sets of components may be operated with different electrical power characteristics (e.g. voltage, clock frequency, etc.). The processor may include multiple domains, where each domain is a collection of hardware and/or logic that operates at the same voltage and frequency point. As used herein, the term “core domain” refers to a domain including the cores of the processor. Further, as used herein, the term “uncore domain” refers to a domain including non-core circuitry such as the mesh network or fabric between the cores. Furthermore, as used herein, the term “memory domain” refers to a domain including memory components of the processor. Although many implementations of a multi-domain processor can be formed on a single semiconductor die, other implementations can be realized by a multi-chip package in which different domains can be present on different semiconductor die of a single package.
- (38) In most processors, the total amount of electrical power that is available (referred to herein as

a “power budget”) in the processor may be limited by one or more maximum thresholds. For example, a processor may have a thermal power limit that is based on the maximum amount of heat that can be removed from the processor during operation. The overall power budget may be allocated by providing different clock frequencies to the different domains of the processor (referred to herein as “frequency budget distribution,” or alternatively as “budget distribution”), such that the total power consumption of the domains is less than or equal to the power budget. However, in some examples, the frequency budget distribution in use may not be suitable for the workload of the different domains at a given point in time. For example, some workloads may require more processing time but less data access, while the reverse situation may apply to other workloads. In such examples, the workload capacity of one or more domains may become constrained, but without constraining the remaining domains. For example, the workload may become constrained by the core domain (referred to herein as “core bound”), by the uncore domain (referred to herein as “uncore bound”), or by the memory domain (referred to herein as “memory bound”). In such situations, it may be difficult to determine a frequency budget distribution that provides relatively high performance per unit of power consumed.

(39) Moreover, in some examples, the types of workloads may vary across different cores of the processor. For examples, different cores may be allocated for use by different entities (e.g., clients) for different tasks, and therefore the cores may have different workload characteristics.

Accordingly, the frequency budget distribution that selected for the entire processor may not be optimal for the different types of workloads (i.e., heterogeneous workloads). Furthermore, the process of changing the frequency budget distribution may require a significant amount of time, and therefore may not be performed quickly enough to respond to changes in dynamic workloads that could otherwise benefit from changes to the frequency distribution. Therefore, such situations may result in missed opportunities to improve the performance of the processor.

(40) In accordance with one or more embodiments, a processor may include circuitry to determine a frequency budget distribution based on current workloads in the processor. The circuitry may receive scalability hint values from each processing core in the processor. The scalability hint values may be generated by a prediction algorithm based on multiple metrics that indicate the current workload of the respective core. The circuitry may determine total core gain and total uncore gain values based at least in part on the scalability hint values received from the cores. The total core gain and total uncore gain values may be used to determine whether the processor is core bound, uncore bound, or memory bound. The circuitry may determine the frequency budget distribution based at least in part on total core gain and total uncore gain values. In this manner, some embodiments may provide a frequency budget distribution that improves the performance of heterogeneous workloads of the cores. Further, some embodiments may allow the frequency budget distribution to be determined with a relatively short response time, and may therefore allow the processor to take advantage of dynamic workloads of short duration. Various details of some embodiments are described further below with reference to FIGS. 26-35. Further, exemplary systems and architectures are described below with reference to FIGS. 1-25.

(41) Exemplary Systems and Architectures

(42) Although the following embodiments are described with reference to particular implementations, embodiments are not limited in this regard. In particular, it is contemplated that similar techniques and teachings of embodiments described herein may be applied to other types of circuits, semiconductor devices, processors, systems, etc. For example, the disclosed embodiments may be implemented in any type of computer system, including server computers (e.g., tower, rack, blade, micro-server and so forth), communications systems, storage systems, desktop computers of any configuration, laptop, notebook, and tablet computers (including 2:1 tablets, phablets and so forth).

(43) In addition, disclosed embodiments can also be used in other devices, such as handheld devices, systems on chip (SoCs), and embedded applications. Some examples of handheld devices

include cellular phones such as smartphones, Internet protocol devices, digital cameras, personal digital assistants (PDAs), and handheld PCs. Embedded applications may typically include a microcontroller, a digital signal processor (DSP), network computers (NetPC), set-top boxes, network hubs, wide area network (WAN) switches, wearable devices, or any other system that can perform the functions and operations taught below. Further, embodiments may be implemented in mobile terminals having standard voice functionality such as mobile phones, smartphones and phablets, and/or in non-mobile terminals without a standard wireless voice function communication capability, such as many wearables, tablets, notebooks, desktops, micro-servers, servers and so forth.

(44) Referring now to FIG. 1, shown is a block diagram of a portion of a system in accordance with an embodiment of the present invention. As shown in FIG. 1, system **100** may include various components, including a processor **110** which as shown is a multicore processor. Processor **110** may be coupled to a power supply **150** via an external voltage regulator **160**, which may perform a first voltage conversion to provide a primary regulated voltage V_{reg} to processor **110**.

(45) As seen, processor **110** may be a single die processor including multiple cores **120a-120n**. In addition, each core may be associated with an integrated voltage regulator (IVR) **125a-125n** which receives the primary regulated voltage and generates an operating voltage to be provided to one or more agents of the processor associated with the IVR. Accordingly, an IVR implementation may be provided to allow for fine-grained control of voltage and thus power and performance of each individual core. As such, each core can operate at an independent voltage and frequency, enabling great flexibility and affording wide opportunities for balancing power consumption with performance. In some embodiments, the use of multiple IVRs enables the grouping of components into separate power planes, such that power is regulated and supplied by the IVR to only those components in the group. During power management, a given power plane of one IVR may be powered down or off when the processor is placed into a certain low power state, while another power plane of another IVR remains active, or fully powered. Similarly, cores **120** may include or be associated with independent clock generation circuitry such as one or more phase lock loops (PLLs) to control operating frequency of each core **120** independently.

(46) Still referring to FIG. 1, additional components may be present within the processor including an input/output interface (IF) **132**, another interface **134**, and an integrated memory controller (IMC) **136**. As seen, each of these components may be powered by another integrated voltage regulator **125.sub.x**. In one embodiment, interface **132** may enable operation for an Intel® Quick Path Interconnect (QPI) interconnect, which provides for point-to-point (PtP) links in a cache coherent protocol that includes multiple layers including a physical layer, a link layer and a protocol layer. In turn, interface **134** may communicate via a Peripheral Component Interconnect Express (PCIe®) protocol.

(47) Also shown is a power control unit (PCU) **138**, which may include circuitry including hardware, software and/or firmware to perform power management operations with regard to processor **110**. As seen, PCU **138** provides control information to external voltage regulator **160** via a digital interface **162** to cause the voltage regulator to generate the appropriate regulated voltage. PCU **138** also provides control information to IVRs **125** via another digital interface **163** to control the operating voltage generated (or to cause a corresponding IVR to be disabled in a low power mode). In various embodiments, PCU **138** may include a variety of power management logic units to perform hardware-based power management. Such power management may be wholly processor controlled (e.g., by various processor hardware, and which may be triggered by workload and/or power, thermal or other processor constraints) and/or the power management may be performed responsive to external sources (such as a platform or power management source or system software).

(48) In FIG. 1, PCU **138** is illustrated as being present as a separate logic of the processor. In other cases, PCU **138** may execute on a given one or more of cores **120**. In some cases, PCU **138** may be

implemented as a microcontroller (dedicated or general-purpose) or other control logic configured to execute its own dedicated power management code, sometimes referred to as P-code. In yet other embodiments, power management operations to be performed by PCU **138** may be implemented externally to a processor, such as by way of a separate power management integrated circuit (PMIC) or another component external to the processor. In yet other embodiments, power management operations to be performed by PCU **138** may be implemented within BIOS or other system software.

(49) Embodiments may be particularly suitable for a multicore processor in which each of multiple cores can operate at an independent voltage and frequency point. As used herein the term “domain” is used to mean a collection of hardware and/or logic that operates at the same voltage and frequency point. In addition, a multicore processor can further include other non-core processing engines such as fixed function units, graphics engines, and so forth. Such processor can include independent domains other than the cores, such as one or more domains associated with a graphics engine (referred to herein as a graphics domain) and one or more domains associated with non-core circuitry, referred to herein as a system agent. Although many implementations of a multi-domain processor can be formed on a single semiconductor die, other implementations can be realized by a multi-chip package in which different domains can be present on different semiconductor die of a single package.

(50) While not shown for ease of illustration, understand that additional components may be present within processor **110** such as non-core logic, and other components such as internal memories, e.g., one or more levels of a cache memory hierarchy and so forth. Furthermore, while shown in the implementation of FIG. **1** with an integrated voltage regulator, embodiments are not so limited. For example, other regulated voltages may be provided to on-chip resources from external voltage regulator **160** or one or more additional external sources of regulated voltages.

(51) Note that the power management techniques described herein may be independent of and complementary to an operating system (OS)-based power management (OSPM) mechanism. According to one example OSPM technique, a processor can operate at various performance states or levels, so-called P-states, namely from P0 to PN. In general, the P1 performance state may correspond to the highest guaranteed performance state that can be requested by an OS. In addition to this P1 state, the OS can further request a higher performance state, namely a P0 state. This P0 state may thus be an opportunistic, overclocking, or turbo mode state in which, when power and/or thermal budget is available, processor hardware can configure the processor or at least portions thereof to operate at a higher than guaranteed frequency. In many implementations, a processor can include multiple so-called bin frequencies above the P1 guaranteed maximum frequency, exceeding to a maximum peak frequency of the particular processor, as fused or otherwise written into the processor during manufacture. In addition, according to one OSPM mechanism, a processor can operate at various power states or levels. With regard to power states, an OSPM mechanism may specify different power consumption states, generally referred to as C-states, C0, C1 to Cn states. When a core is active, it runs at a C0 state, and when the core is idle it may be placed in a core low power state, also called a core non-zero C-state (e.g., C1-C6 states), with each C-state being at a lower power consumption level (such that C6 is a deeper low power state than C1, and so forth).

(52) Understand that many different types of power management techniques may be used individually or in combination in different embodiments. As representative examples, a power controller may control the processor to be power managed by some form of dynamic voltage frequency scaling (DVFS) in which an operating voltage and/or operating frequency of one or more cores or other processor logic may be dynamically controlled to reduce power consumption in certain situations. In an example, DVFS may be performed using Enhanced Intel SpeedStep™ technology available from Intel Corporation, Santa Clara, CA, to provide optimal performance at a lowest power consumption level. In another example, DVFS may be performed using Intel TurboBoost™ technology to enable one or more cores or other compute engines to operate at a

higher than guaranteed operating frequency based on conditions (e.g., workload and availability).

(53) Another power management technique that may be used in certain examples is dynamic swapping of workloads between different compute engines. For example, the processor may include asymmetric cores or other processing engines that operate at different power consumption levels, such that in a power constrained situation, one or more workloads can be dynamically switched to execute on a lower power core or other compute engine. Another exemplary power management technique is hardware duty cycling (HDC), which may cause cores and/or other compute engines to be periodically enabled and disabled according to a duty cycle, such that one or more cores may be made inactive during an inactive period of the duty cycle and made active during an active period of the duty cycle.

(54) Power management techniques also may be used when constraints exist in an operating environment. For example, when a power and/or thermal constraint is encountered, power may be reduced by reducing operating frequency and/or voltage. Other power management techniques include throttling instruction execution rate or limiting scheduling of instructions. Still further, it is possible for instructions of a given instruction set architecture to include express or implicit direction as to power management operations. Although described with these particular examples, understand that many other power management techniques may be used in particular embodiments.

(55) Embodiments can be implemented in processors for various markets including server processors, desktop processors, mobile processors and so forth. Referring now to FIG. 2, shown is a block diagram of a processor in accordance with an embodiment of the present invention. As shown in FIG. 2, processor **200** may be a multicore processor including a plurality of cores **210a-210n**. In one embodiment, each such core may be of an independent power domain and can be configured to enter and exit active states and/or maximum performance states based on workload. One or more cores **210** may be heterogeneous to the other cores, e.g., having different micro-architectures, instruction set architectures, pipeline depths, power and performance capabilities. The various cores may be coupled via an interconnect **215** to a system agent **220** that includes various components. As seen, the system agent **220** may include a shared cache **230** which may be a last level cache. In addition, the system agent may include an integrated memory controller **240** to communicate with a system memory (not shown in FIG. 2), e.g., via a memory bus. The system agent **220** also includes various interfaces **250** and a power control unit **255**, which may include logic to perform the power management techniques described herein.

(56) In addition, by interfaces **250a-250n**, connection can be made to various off-chip components such as peripheral devices, mass storage and so forth. While shown with this particular implementation in the embodiment of FIG. 2, the scope of the present invention is not limited in this regard.

(57) Referring now to FIG. 3, shown is a block diagram of a multi-domain processor in accordance with another embodiment of the present invention. As shown in the embodiment of FIG. 3, processor **300** includes multiple domains. Specifically, a core domain **310** can include a plurality of cores **310a-310n**, a graphics domain **320** can include one or more graphics engines, and a system agent domain **350** may further be present. In some embodiments, system agent domain **350** may execute at an independent frequency than the core domain and may remain powered on at all times to handle power control events and power management such that domains **310** and **320** can be controlled to dynamically enter into and exit high power and low power states. Each of domains **310** and **320** may operate at different voltage and/or power. Note that while only shown with three domains, understand the scope of the present invention is not limited in this regard and additional domains can be present in other embodiments. For example, multiple core domains may be present each including at least one core.

(58) In general, each of the cores **310a-310n** may further include low level caches in addition to various execution units and additional processing elements. In turn, the various cores may be coupled to each other and to a shared cache memory formed of a plurality of units of a last level

cache (LLC) **340a-340n**. In various embodiments, LLC **340** may be shared amongst the cores and the graphics engine, as well as various media processing circuitry. As seen, a ring interconnect **330** thus couples the cores together, and provides interconnection between the cores, graphics domain **320** and system agent domain **350**. In one embodiment, interconnect **330** can be part of the core domain. However, in other embodiments the ring interconnect can be of its own domain.

(59) As further seen, system agent domain **350** may include display controller **352** which may provide control of and an interface to an associated display. As further seen, system agent domain **350** may include a power control unit **355** which can include logic to perform the power management techniques described herein.

(60) As further seen in FIG. 3, processor **300** can further include an integrated memory controller (IMC) **370** that can provide for an interface to a system memory, such as a dynamic random access memory (DRAM). Multiple interfaces **380a-380n** may be present to enable interconnection between the processor and other circuitry. For example, in one embodiment at least one direct media interface (DMI) interface may be provided as well as one or more PCIe™ interfaces. Still further, to provide for communications between other agents such as additional processors or other circuitry, one or more QPI interfaces may also be provided. Although shown at this high level in the embodiment of FIG. 3, understand the scope of the present invention is not limited in this regard.

(61) Referring to FIG. 4, an embodiment of a processor including multiple cores is illustrated. Processor **400** includes any processor or processing device, such as a microprocessor, an embedded processor, a digital signal processor (DSP), a network processor, a handheld processor, an application processor, a co-processor, a system on a chip (SoC), or other device to execute code. Processor **400**, in one embodiment, includes at least two cores—cores **401** and **402**, which may include asymmetric cores or symmetric cores (the illustrated embodiment). However, processor **400** may include any number of processing elements that may be symmetric or asymmetric.

(62) In one embodiment, a processing element refers to hardware or logic to support a software thread. Examples of hardware processing elements include: a thread unit, a thread slot, a thread, a process unit, a context, a context unit, a logical processor, a hardware thread, a core, and/or any other element, which is capable of holding a state for a processor, such as an execution state or architectural state. In other words, a processing element, in one embodiment, refers to any hardware capable of being independently associated with code, such as a software thread, operating system, application, or other code. A physical processor typically refers to an integrated circuit, which potentially includes any number of other processing elements, such as cores or hardware threads.

(63) A core often refers to logic located on an integrated circuit capable of maintaining an independent architectural state, wherein each independently maintained architectural state is associated with at least some dedicated execution resources. In contrast to cores, a hardware thread typically refers to any logic located on an integrated circuit capable of maintaining an independent architectural state, wherein the independently maintained architectural states share access to execution resources. As can be seen, when certain resources are shared and others are dedicated to an architectural state, the line between the nomenclature of a hardware thread and core overlaps. Yet often, a core and a hardware thread are viewed by an operating system as individual logical processors, where the operating system is able to individually schedule operations on each logical processor.

(64) Physical processor **400**, as illustrated in FIG. 4, includes two cores, cores **401** and **402**. Here, cores **401** and **402** are considered symmetric cores, i.e., cores with the same configurations, functional units, and/or logic. In another embodiment, core **401** includes an out-of-order processor core, while core **402** includes an in-order processor core. However, cores **401** and **402** may be individually selected from any type of core, such as a native core, a software managed core, a core adapted to execute a native instruction set architecture (ISA), a core adapted to execute a translated

ISA, a co-designed core, or other known core. Yet to further the discussion, the functional units illustrated in core **401** are described in further detail below, as the units in core **402** operate in a similar manner.

(65) As depicted, core **401** includes two architectural state registers **401a** and **401b**, which may be associated with two hardware threads (also referred to as hardware thread slots). Therefore, software entities, such as an operating system, in one embodiment potentially view processor **400** as four separate processors, i.e., four logical processors or processing elements capable of executing four software threads concurrently. As alluded to above, a first thread is associated with architecture state registers **401a**, a second thread is associated with architecture state registers **401b**, a third thread may be associated with architecture state registers **402a**, and a fourth thread may be associated with architecture state registers **402b**. Here, the architecture state registers (**401a**, **401b**, **402a**, and **402b**) may be associated with processing elements, thread slots, or thread units, as described above. As illustrated, architecture state registers **401a** are replicated in architecture state registers **401b**, so individual architecture states/contexts are capable of being stored for logical processor **401a** and logical processor **401b**. In core **401**, other smaller resources, such as instruction pointers and renaming logic in allocator and renamer block **430** may also be replicated for threads **401a** and **401b**. Some resources, such as re-order buffers in reorder/retirement unit **435**, branch target buffer and instruction translation lookaside buffer (BTB and I-TLB) **420**, load/store buffers, and queues may be shared through partitioning. Other resources, such as general purpose internal registers, page-table base register(s), low-level data-cache and data-TLB **450**, execution unit(s) **440**, and portions of reorder/retirement unit **435** are potentially fully shared.

(66) Processor **400** often includes other resources, which may be fully shared, shared through partitioning, or dedicated by/to processing elements. In FIG. 4, an embodiment of a purely exemplary processor with illustrative logical units/resources of a processor is illustrated. Note that a processor may include, or omit, any of these functional units, as well as include any other known functional units, logic, or firmware not depicted. As illustrated, core **401** includes a simplified, representative out-of-order (OOO) processor core. But an in-order processor may be utilized in different embodiments.

(67) Core **401** further includes decode module **425** coupled to a fetch unit to decode fetched elements. Fetch logic, in one embodiment, includes individual sequencers associated with thread slots **401a**, **401b**, respectively. Usually core **401** is associated with a first ISA, which defines/specifies instructions executable on processor **400**. Often machine code instructions that are part of the first ISA include a portion of the instruction (referred to as an opcode), which references/specifies an instruction or operation to be performed. Decode module **425** includes circuitry that recognizes these instructions from their opcodes and passes the decoded instructions on in the pipeline for processing as defined by the first ISA. For example, decoder module **425**, in one embodiment, includes logic designed or adapted to recognize specific instructions, such as transactional instructions. As a result of the recognition by the decoder module **425**, the architecture or core **401** takes specific, predefined actions to perform tasks associated with the appropriate instruction. It is important to note that any of the tasks, blocks, operations, and methods described herein may be performed in response to a single or multiple instructions; some of which may be new or old instructions.

(68) In one example, allocator and renamer block **430** includes an allocator to reserve resources, such as register files to store instruction processing results. However, threads **401a** and **401b** are potentially capable of out-of-order execution, where allocator and renamer block **430** also reserves other resources, such as reorder buffers to track instruction results. The renamer block **430** may also include a register renamer to rename program/instruction reference registers to other registers internal to processor **400**. Reorder/retirement unit **435** includes components, such as the reorder buffers mentioned above, load buffers, and store buffers, to support out-of-order execution and later in-order retirement of instructions executed out-of-order.

(69) Scheduler and execution unit(s) block **440**, in one embodiment, includes a scheduler unit to schedule instructions/operation on execution units. For example, a floating point instruction is scheduled on a port of an execution unit that has an available floating point execution unit. Register files associated with the execution units are also included to store information instruction processing results. Exemplary execution units include a floating point execution unit, an integer execution unit, a jump execution unit, a load execution unit, a store execution unit, and other known execution units.

(70) Lower level data cache and data translation lookaside buffer (D-TLB) **450** are coupled to execution unit(s) **440**. The data cache is to store recently used/operated on elements, such as data operands, which are potentially held in memory coherency states. The D-TLB is to store recent virtual/linear to physical address translations. As a specific example, a processor may include a page table structure to break physical memory into a plurality of virtual pages.

(71) Here, cores **401** and **402** share access to higher-level or further-out cache **410**, which is to cache recently fetched elements. Note that higher-level or further-out refers to cache levels increasing or getting further away from the execution unit(s). In one embodiment, higher-level cache **410** is a last-level data cache—last cache in the memory hierarchy on processor **400**—such as a second or third level data cache. However, higher level cache **410** is not so limited, as it may be associated with or includes an instruction cache. A trace cache—a type of instruction cache—instead may be coupled after decoder module **425** to store recently decoded traces.

(72) In the depicted configuration, processor **400** also includes bus interface **405** and a power control unit **460**, which may perform power management in accordance with an embodiment of the present invention. In this scenario, bus interface **405** is to communicate with devices external to processor **400**, such as system memory and other components.

(73) A memory controller **470** may interface with other devices such as one or many memories. In an example, bus interface **405** includes a ring interconnect with a memory controller for interfacing with a memory and a graphics controller for interfacing with a graphics processor. In an SoC environment, even more devices, such as a network interface, coprocessors, memory, graphics processor, and any other known computer devices/interface may be integrated on a single die or integrated circuit to provide small form factor with high functionality and low power consumption.

(74) Referring now to FIG. 5, shown is a block diagram of a micro-architecture of a processor core in accordance with one embodiment of the present invention. As shown in FIG. 5, processor core **500** may be a multi-stage pipelined out-of-order processor. Core **500** may operate at various voltages based on a received operating voltage, which may be received from an integrated voltage regulator or external voltage regulator.

(75) As seen in FIG. 5, core **500** includes front end units **510**, which may be used to fetch instructions to be executed and prepare them for use later in the processor pipeline. For example, front end units **510** may include a fetch unit **501**, an instruction cache **503**, and an instruction decoder **505**. In some implementations, front end units **510** may further include a trace cache, along with microcode storage as well as a micro-operation storage.

(76) Fetch unit **501** may fetch macro-instructions, e.g., from memory or instruction cache **503**, and feed them to instruction decoder **505** to decode them into primitives, i.e., micro-operations for execution by the processor.

(77) Coupled between front end units **510** and execution units **520** is an out-of-order (OOO) engine **515** that may be used to receive the micro-instructions and prepare them for execution. More specifically OOO engine **515** may include various buffers to re-order micro-instruction flow and allocate various resources needed for execution, as well as to provide renaming of logical registers onto storage locations within various register files such as register file **530** and extended register file **535**. Register file **530** may include separate register files for integer and floating point operations. For purposes of configuration, control, and additional operations, a set of machine specific registers (MSRs) **538** may also be present and accessible to various logic within core **500**.

(and external to the core).

(78) Various resources may be present in execution units **520**, including, for example, various integer, floating point, and single instruction multiple data (SIMD) logic units, among other specialized hardware. For example, such execution units may include one or more arithmetic logic units (ALUs) **522** and one or more vector execution units **524**, among other such execution units.

(79) Results from the execution units may be provided to retirement logic, namely a reorder buffer (ROB) **540**. More specifically, ROB **540** may include various arrays and logic to receive information associated with instructions that are executed. This information is then examined by ROB **540** to determine whether the instructions can be validly retired and result data committed to the architectural state of the processor, or whether one or more exceptions occurred that prevent a proper retirement of the instructions. Of course, ROB **540** may handle other operations associated with retirement.

(80) As shown in FIG. 5, ROB **540** is coupled to a cache **550** which, in one embodiment may be a low level cache (e.g., an L1 cache) although the scope of the present invention is not limited in this regard. Also, execution units **520** can be directly coupled to cache **550**. From cache **550**, data communication may occur with higher level caches, system memory and so forth. While shown with this high level in the embodiment of FIG. 5, understand the scope of the present invention is not limited in this regard. For example, while the implementation of FIG. 5 is with regard to an out-of-order machine such as of an Intel® x86 instruction set architecture (ISA), the scope of the present invention is not limited in this regard. That is, other embodiments may be implemented in an in-order processor, a reduced instruction set computing (RISC) processor such as an ARM-based processor, or a processor of another type of ISA that can emulate instructions and operations of a different ISA via an emulation engine and associated logic circuitry.

(81) Referring now to FIG. 6, shown is a block diagram of a micro-architecture of a processor core in accordance with another embodiment. In the embodiment of FIG. 6, core **600** may be a low power core of a different micro-architecture, such as an Intel® Atom™-based processor having a relatively limited pipeline depth designed to reduce power consumption. As seen, core **600** includes an instruction cache **610** coupled to provide instructions to an instruction decoder **615**. A branch predictor **605** may be coupled to instruction cache **610**. Note that instruction cache **610** may further be coupled to another level of a cache memory, such as an L2 cache (not shown for ease of illustration in FIG. 6). In turn, instruction decoder **615** provides decoded instructions to an issue queue (IQ) **620** for storage and delivery to a given execution pipeline. A microcode ROM **618** is coupled to instruction decoder **615**.

(82) A floating point pipeline **630** includes a floating point (FP) register file **632** which may include a plurality of architectural registers of a given bit width such as 128, 256 or 512 bits. Pipeline **630** includes a floating point scheduler **634** to schedule instructions for execution on one of multiple execution units of the pipeline. In the embodiment shown, such execution units include an arithmetic logic unit (ALU) **635**, a shuffle unit **636**, and a floating point (FP) adder **638**. In turn, results generated in these execution units may be provided back to buffers and/or registers of register file **632**. Of course understand while shown with these few example execution units, additional or different floating point execution units may be present in another embodiment.

(83) An integer pipeline **640** also may be provided. In the embodiment shown, pipeline **640** includes an integer (INT) register file **642** which may include a plurality of architectural registers of a given bit width such as 128 or 256 bits. Pipeline **640** includes an integer execution (IE) scheduler **644** to schedule instructions for execution on one of multiple execution units of the pipeline. In the embodiment shown, such execution units include an ALU **645**, a shifter unit **646**, and a jump execution unit (JEU) **648**. In turn, results generated in these execution units may be provided back to buffers and/or registers of register file **642**. Of course, understand while shown with these few example execution units, additional or different integer execution units may be present in another embodiment.

(84) A memory execution (ME) scheduler **650** may schedule memory operations for execution in an address generation unit (AGU) **652**, which is also coupled to a TLB **654**. As seen, these structures may couple to a data cache **660**, which may be a L0 and/or L1 data cache that in turn couples to additional levels of a cache memory hierarchy, including an L2 cache memory.

(85) To provide support for out-of-order execution, an allocator/renamer **670** may be provided, in addition to a reorder buffer **680**, which is configured to reorder instructions executed out of order for retirement in order. Although shown with this particular pipeline architecture in the illustration of FIG. 6, understand that many variations and alternatives are possible.

(86) Note that in a processor having asymmetric cores, such as in accordance with the micro-architectures of FIGS. 5 and 6, workloads may be dynamically swapped between the cores for power management reasons, as these cores, although having different pipeline designs and depths, may be of the same or related ISA. Such dynamic core swapping may be performed in a manner transparent to a user application (and possibly kernel also).

(87) Referring to FIG. 7, shown is a block diagram of a micro-architecture of a processor core in accordance with yet another embodiment. As illustrated in FIG. 7, a core **700** may include a multi-staged in-order pipeline to execute at very low power consumption levels. As one such example, core **700** may have a micro-architecture in accordance with an ARM Cortex A53 design available from ARM Holdings, LTD., Sunnyvale, CA. In an implementation, an 8-stage pipeline may be provided that is configured to execute both 32-bit and 64-bit code. Core **700** includes a fetch unit **710** that is configured to fetch instructions and provide them to a decode unit **715**, which may decode the instructions, e.g., macro-instructions of a given ISA such as an ARMv8 ISA. Note further that a queue **730** may couple to decode unit **715** to store decoded instructions. Decoded instructions are provided to an issue logic **725**, where the decoded instructions may be issued to a given one of multiple execution units.

(88) With further reference to FIG. 7, issue logic **725** may issue instructions to one of multiple execution units. In the embodiment shown, these execution units include an integer unit **735**, a multiply unit **740**, a floating point/vector unit **750**, a dual issue unit **760**, and a load/store unit **770**. The results of these different execution units may be provided to a writeback (WB) unit **780**. Understand that while a single writeback unit is shown for ease of illustration, in some implementations separate writeback units may be associated with each of the execution units. Furthermore, understand that while each of the units and logic shown in FIG. 7 is represented at a high level, a particular implementation may include more or different structures. A processor designed using one or more cores having a pipeline as in FIG. 7 may be implemented in many different end products, extending from mobile devices to server systems.

(89) Referring to FIG. 8, shown is a block diagram of a micro-architecture of a processor core in accordance with a still further embodiment. As illustrated in FIG. 8, a core **800** may include a multi-stage multi-issue out-of-order pipeline to execute at very high performance levels (which may occur at higher power consumption levels than core **700** of FIG. 7). As one such example, processor **800** may have a microarchitecture in accordance with an ARM Cortex A57 design. In an implementation, a 15 (or greater)-stage pipeline may be provided that is configured to execute both 32-bit and 64-bit code. In addition, the pipeline may provide for 3 (or greater)-wide and 3 (or greater)-issue operation. Core **800** includes a fetch unit **810** that is configured to fetch instructions and provide them to a decoder/renamer/dispatcher unit **815** coupled to a cache **820**. Unit **815** may decode the instructions, e.g., macro-instructions of an ARMv8 instruction set architecture, rename register references within the instructions, and dispatch the instructions (eventually) to a selected execution unit. Decoded instructions may be stored in a queue **825**. Note that while a single queue structure is shown for ease of illustration in FIG. 8, understand that separate queues may be provided for each of the multiple different types of execution units.

(90) Also shown in FIG. 8 is an issue logic **830** from which decoded instructions stored in queue **825** may be issued to a selected execution unit. Issue logic **830** also may be implemented in a

particular embodiment with a separate issue logic for each of the multiple different types of execution units to which issue logic **830** couples.

(91) Decoded instructions may be issued to a given one of multiple execution units. In the embodiment shown, these execution units include one or more integer units **835**, a multiply unit **840**, a floating point/vector unit **850**, a branch unit **860**, and a load/store unit **870**. In an embodiment, floating point/vector unit **850** may be configured to handle SIMD or vector data of 128 or 256 bits. Still further, floating point/vector execution unit **850** may perform IEEE-754 double precision floating-point operations. The results of these different execution units may be provided to a writeback unit **880**. Note that in some implementations separate writeback units may be associated with each of the execution units. Furthermore, understand that while each of the units and logic shown in FIG. **8** is represented at a high level, a particular implementation may include more or different structures.

(92) Note that in a processor having asymmetric cores, such as in accordance with the micro-architectures of FIGS. **7** and **8**, workloads may be dynamically swapped for power management reasons, as these cores, although having different pipeline designs and depths, may be of the same or related ISA. Such dynamic core swapping may be performed in a manner transparent to a user application (and possibly kernel also).

(93) A processor designed using one or more cores having pipelines as in any one or more of FIGS. **5-8** may be implemented in many different end products, extending from mobile devices to server systems. Referring now to FIG. **9**, shown is a block diagram of a processor in accordance with another embodiment of the present invention. In the embodiment of FIG. **9**, processor **900** may be a SoC including multiple domains, each of which may be controlled to operate at an independent operating voltage and operating frequency. As a specific illustrative example, processor **900** may be an Intel® Architecture Core™-based processor such as an i3, i5, i7 or another such processor available from Intel Corporation. However, other low power processors such as available from Advanced Micro Devices, Inc. (AMD) of Sunnyvale, CA, an ARM-based design from ARM Holdings, Ltd. or licensee thereof or a MIPS-based design from MIPS Technologies, Inc. of Sunnyvale, CA, or their licensees or adopters may instead be present in other embodiments such as an Apple A7 processor, a Qualcomm Snapdragon processor, or Texas Instruments OMAP processor. Such SoC may be used in a low power system such as a smartphone, tablet computer, phablet computer, Ultrabook™ computer or other portable computing device, which may incorporate a heterogeneous system architecture having a heterogeneous system architecture-based processor design.

(94) In the high level view shown in FIG. **9**, processor **900** includes a plurality of core units **910a-910n**. Each core unit may include one or more processor cores, one or more cache memories and other circuitry. Each core unit **910** may support one or more instruction sets (e.g., an x86 instruction set (with some extensions that have been added with newer versions); a MIPS instruction set; an ARM instruction set (with optional additional extensions such as NEON)) or other instruction set or combinations thereof. Note that some of the core units may be heterogeneous resources (e.g., of a different design). In addition, each such core may be coupled to a cache memory (not shown) which in an embodiment may be a shared level two (L2) cache memory. A non-volatile storage **930** may be used to store various program and other data. For example, this storage may be used to store at least portions of microcode, boot information such as a BIOS, other system software or so forth.

(95) Each core unit **910** may also include an interface such as a bus interface unit to enable interconnection to additional circuitry of the processor. In an embodiment, each core unit **910** couples to a coherent fabric that may act as a primary cache coherent on-die interconnect that in turn couples to a memory controller **935**. In turn, memory controller **935** controls communications with a memory such as a DRAM (not shown for ease of illustration in FIG. **9**).

(96) In addition to core units, additional processing engines are present within the processor,

including at least one graphics unit **920** which may include one or more graphics processing units (GPUs) to perform graphics processing as well as to possibly execute general purpose operations on the graphics processor (so-called GPGPU operation). In addition, at least one image signal processor **925** may be present. Signal processor **925** may be configured to process incoming image data received from one or more capture devices, either internal to the SoC or off-chip.

(97) Other accelerators also may be present. In the illustration of FIG. **9**, a video coder **950** may perform coding operations including encoding and decoding for video information, e.g., providing hardware acceleration support for high definition video content. A display controller **955** further may be provided to accelerate display operations including providing support for internal and external displays of a system. In addition, a security processor **945** may be present to perform security operations such as secure boot operations, various cryptography operations and so forth.

(98) Each of the units may have its power consumption controlled via a power manager **940**, which may include control logic to perform the various power management techniques described herein.

(99) In some embodiments, processor **900** may further include a non-coherent fabric coupled to the coherent fabric to which various peripheral devices may couple. One or more interfaces **960a-960d** enable communication with one or more off-chip devices. Such communications may be via a variety of communication protocols such as PCIe™ GPIO, USB, I.sup.2C, UART, MIPI, SDIO, DDR, SPI, HDMI, among other types of communication protocols. Although shown at this high level in the embodiment of FIG. **9**, understand the scope of the present invention is not limited in this regard.

(100) Referring now to FIG. **10**, shown is a block diagram of a representative SoC. In the embodiment shown, SoC **1000** may be a multi-core SoC configured for low power operation to be optimized for incorporation into a smartphone or other low power device such as a tablet computer or other portable computing device. As an example, SoC **1000** may be implemented using asymmetric or different types of cores, such as combinations of higher power and/or low power cores, e.g., out-of-order cores and in-order cores. In different embodiments, these cores may be based on an Intel® Architecture™ core design or an ARM architecture design. In yet other embodiments, a mix of Intel and ARM cores may be implemented in a given SoC.

(101) As seen in FIG. **10**, SoC **1000** includes a first core domain **1010** having a plurality of first cores **1012a-1012d**. In an example, these cores may be low power cores such as in-order cores. In one embodiment, these first cores may be implemented as ARM Cortex A53 cores. In turn, these cores couple to a cache memory **1015** of core domain **1010**. In addition, SoC **1000** includes a second core domain **1020**. In the illustration of FIG. **10**, second core domain **1020** has a plurality of second cores **1022a-1022d**. In an example, these cores may be higher power-consuming cores than first cores **1012**. In an embodiment, the second cores may be out-of-order cores, which may be implemented as ARM Cortex A57 cores. In turn, these cores couple to a cache memory **1025** of core domain **1020**. Note that while the example shown in FIG. **10** includes **4** cores in each domain, understand that more or fewer cores may be present in a given domain in other examples.

(102) With further reference to FIG. **10**, a graphics domain **1030** also is provided, which may include one or more graphics processing units (GPUs) configured to independently execute graphics workloads, e.g., provided by one or more cores of core domains **1010** and **1020**. As an example, GPU domain **1030** may be used to provide display support for a variety of screen sizes, in addition to providing graphics and display rendering operations.

(103) As seen, the various domains couple to a coherent interconnect **1040**, which in an embodiment may be a cache coherent interconnect fabric that in turn couples to an integrated memory controller **1050**. Coherent interconnect **1040** may include a shared cache memory, such as an L3 cache, in some examples. In an embodiment, memory controller **1050** may be a direct memory controller to provide for multiple channels of communication with an off-chip memory, such as multiple channels of a DRAM (not shown for ease of illustration in FIG. **10**).

(104) In different examples, the number of the core domains may vary. For example, for a low

power SoC suitable for incorporation into a mobile computing device, a limited number of core domains such as shown in FIG. **10** may be present. Still further, in such low power SoCs, core domain **1020** including higher power cores may have fewer numbers of such cores. For example, in one implementation two cores **1022** may be provided to enable operation at reduced power consumption levels. In addition, the different core domains may also be coupled to an interrupt controller to enable dynamic swapping of workloads between the different domains.

(105) In yet other embodiments, a greater number of core domains, as well as additional optional IP logic may be present, in that an SoC can be scaled to higher performance (and power) levels for incorporation into other computing devices, such as desktops, servers, high performance computing systems, base stations forth. As one such example, 4 core domains each having a given number of out-of-order cores may be provided. Still further, in addition to optional GPU support (which as an example may take the form of a GPGPU), one or more accelerators to provide optimized hardware support for particular functions (e.g. web serving, network processing, switching or so forth) also may be provided. In addition, an input/output interface may be present to couple such accelerators to off-chip components.

(106) Referring now to FIG. **11**, shown is a block diagram of another example SoC. In the embodiment of FIG. **11**, SoC **1100** may include various circuitry to enable high performance for multimedia applications, communications and other functions. As such, SoC **1100** is suitable for incorporation into a wide variety of portable and other devices, such as smartphones, tablet computers, smart TVs and so forth. In the example shown, SoC **1100** includes a central processor unit (CPU) domain **1110**. In an embodiment, a plurality of individual processor cores may be present in CPU domain **1110**. As one example, CPU domain **1110** may be a quad core processor having 4 multithreaded cores. Such processors may be homogeneous or heterogeneous processors, e.g., a mix of low power and high power processor cores.

(107) In turn, a GPU domain **1120** is provided to perform advanced graphics processing in one or more GPUs to handle graphics and compute APIs. A DSP unit **1130** may provide one or more low power DSPs for handling low-power multimedia applications such as music playback, audio/video and so forth, in addition to advanced calculations that may occur during execution of multimedia instructions. In turn, a communication unit **1140** may include various components to provide connectivity via various wireless protocols, such as cellular communications (including 3G/4G LTE), wireless local area protocols such as Bluetooth™, IEEE 802.11, and so forth.

(108) Still further, a multimedia processor **1150** may be used to perform capture and playback of high definition video and audio content, including processing of user gestures. A sensor unit **1160** may include a plurality of sensors and/or a sensor controller to interface to various off-chip sensors present in a given platform. An image signal processor (ISP) **1170** may perform image processing with regard to captured content from one or more cameras of a platform, including still and video cameras.

(109) A display processor **1180** may provide support for connection to a high definition display of a given pixel density, including the ability to wirelessly communicate content for playback on such display. Still further, a location unit **1190** may include a Global Positioning System (GPS) receiver with support for multiple GPS constellations to provide applications highly accurate positioning information obtained using as such GPS receiver. Understand that while shown with this particular set of components in the example of FIG. **11**, many variations and alternatives are possible.

(110) Referring now to FIG. **12**, shown is a block diagram of an example system with which embodiments can be used. As seen, system **1200** may be a smartphone or other wireless communicator. A baseband processor **1205** is configured to perform various signal processing with regard to communication signals to be transmitted from or received by the system. In turn, baseband processor **1205** is coupled to an application processor **1210**, which may be a main CPU of the system to execute an OS and other system software, in addition to user applications such as many well-known social media and multimedia apps. Application processor **1210** may further be

configured to perform a variety of other computing operations for the device.

(111) In turn, application processor **1210** can couple to a user interface/display **1220**, e.g., a touch screen display. In addition, application processor **1210** may couple to a memory system including a non-volatile memory, namely a flash memory **1230** and a system memory, namely a dynamic random access memory (DRAM) **1235**. As further seen, application processor **1210** further couples to a capture device **1241** such as one or more image capture devices that can record video and/or still images.

(112) Still referring to FIG. **12**, a universal integrated circuit card (UICC) **1246** comprising a subscriber identity module and possibly a secure storage and cryptoprocessor is also coupled to application processor **1210**. System **1200** may further include a security processor **1250** that may couple to application processor **1210**. A plurality of sensors **1225** may couple to application processor **1210** to enable input of a variety of sensed information such as accelerometer and other environmental information. An audio output device **1295** may provide an interface to output sound, e.g., in the form of voice communications, played or streaming audio data and so forth.

(113) As further illustrated, a near field communication (NFC) contactless interface **1260** is provided that communicates in a NFC near field via an NFC antenna **1265**. While separate antennae are shown in FIG. **12**, understand that in some implementations one antenna or a different set of antennae may be provided to enable various wireless functionality.

(114) A power management integrated circuit (PMIC) **1215** couples to application processor **1210** to perform platform level power management. To this end, PMIC **1215** may issue power management requests to application processor **1210** to enter certain low power states as desired. Furthermore, based on platform constraints, PMIC **1215** may also control the power level of other components of system **1200**.

(115) To enable communications to be transmitted and received, various circuitry may be coupled between baseband processor **1205** and an antenna **1290**. Specifically, a radio frequency (RF) transceiver **1270** and a wireless local area network (WLAN) transceiver **1275** may be present. In general, RF transceiver **1270** may be used to receive and transmit wireless data and calls according to a given wireless communication protocol such as 3G or 4G wireless communication protocol such as in accordance with a code division multiple access (CDMA), global system for mobile communication (GSM), long term evolution (LTE) or other protocol. In addition a GPS sensor **1280** may be present. Other wireless communications such as receipt or transmission of radio signals, e.g., AM/FM and other signals may also be provided. In addition, via WLAN transceiver **1275**, local wireless communications can also be realized.

(116) Referring now to FIG. **13**, shown is a block diagram of another example system with which embodiments may be used. In the illustration of FIG. **13**, system **1300** may be mobile low-power system such as a tablet computer, 2:1 tablet, phablet or other convertible or standalone tablet system. As illustrated, a SoC **1310** is present and may be configured to operate as an application processor for the device.

(117) A variety of devices may couple to SoC **1310**. In the illustration shown, a memory subsystem includes a flash memory **1340** and a DRAM **1345** coupled to SoC **1310**. In addition, a touch panel **1320** is coupled to the SoC **1310** to provide display capability and user input via touch, including provision of a virtual keyboard on a display of touch panel **1320**. To provide wired network connectivity, SoC **1310** couples to an Ethernet interface **1330**. A peripheral hub **1325** is coupled to SoC **1310** to enable interfacing with various peripheral devices, such as may be coupled to system **1300** by any of various ports or other connectors.

(118) In addition to internal power management circuitry and functionality within SoC **1310**, a PMIC **1380** is coupled to SoC **1310** to provide platform-based power management, e.g., based on whether the system is powered by a battery **1390** or AC power via an AC adapter **1395**. In addition to this power source-based power management, PMIC **1380** may further perform platform power management activities based on environmental and usage conditions. Still further, PMIC **1380** may

communicate control and status information to SoC **1310** to cause various power management actions within SoC **1310**.

(119) Still referring to FIG. **13**, to provide for wireless capabilities, a WLAN unit **1350** is coupled to SoC **1310** and in turn to an antenna **1355**. In various implementations, WLAN unit **1350** may provide for communication according to one or more wireless protocols.

(120) As further illustrated, a plurality of sensors **1360** may couple to SoC **1310**. These sensors may include various accelerometer, environmental and other sensors, including user gesture sensors. Finally, an audio codec **1365** is coupled to SoC **1310** to provide an interface to an audio output device **1370**. Of course understand that while shown with this particular implementation in FIG. **13**, many variations and alternatives are possible.

(121) Referring now to FIG. **14**, shown is a block diagram of a representative computer system **1400** such as notebook, Ultrabook™ or other small form factor system. A processor **1410**, in one embodiment, includes a microprocessor, multi-core processor, multithreaded processor, an ultra low voltage processor, an embedded processor, or other known processing element. In the illustrated implementation, processor **1410** acts as a main processing unit and central hub for communication with many of the various components of the system **1400**, and may include power management circuitry as described herein. As one example, processor **1410** is implemented as a SoC.

(122) Processor **1410**, in one embodiment, communicates with a system memory **1415**. As an illustrative example, the system memory **1415** is implemented via multiple memory devices or modules to provide for a given amount of system memory.

(123) To provide for persistent storage of information such as data, applications, one or more operating systems and so forth, a mass storage **1420** may also couple to processor **1410**. In various embodiments, to enable a thinner and lighter system design as well as to improve system responsiveness, this mass storage may be implemented via a SAD or the mass storage may primarily be implemented using a hard disk drive (HDD) with a smaller amount of SAD storage to act as a SAD cache to enable non-volatile storage of context state and other such information during power down events so that a fast power up can occur on re-initiation of system activities. Also shown in FIG. **14**, a flash device **1422** may be coupled to processor **1410**, e.g., via a serial peripheral interface (SPI). This flash device may provide for non-volatile storage of system software, including a basic input/output software (BIOS) as well as other firmware of the system.

(124) Various input/output (I/O) devices may be present within system **1400**. Specifically shown in the embodiment of FIG. **14** is a display **1424** which may be a high definition LCD or LED panel that further provides for a touch screen **1425**. In one embodiment, display **1424** may be coupled to processor **1410** via a display interconnect that can be implemented as a high performance graphics interconnect. Touch screen **1425** may be coupled to processor **1410** via another interconnect, which in an embodiment can be an I.sup.2C interconnect. As further shown in FIG. **14**, in addition to touch screen **1425**, user input by way of touch can also occur via a touch pad **1430** which may be configured within the chassis and may also be coupled to the same I.sup.2C interconnect as touch screen **1425**.

(125) For perceptual computing and other purposes, various sensors may be present within the system and may be coupled to processor **1410** in different manners. Certain inertial and environmental sensors may couple to processor **1410** through a sensor hub **1440**, e.g., via an I.sup.2C interconnect. In the embodiment shown in FIG. **14**, these sensors may include an accelerometer **1441**, an ambient light sensor (ALS) **1442**, a compass **1443** and a gyroscope **1444**. Other environmental sensors may include one or more thermal sensors **1446** which in some embodiments couple to processor **1410** via a system management bus (SMBus) bus.

(126) As also seen in FIG. **14**, various peripheral devices may couple to processor **1410** via a low pin count (LPC) interconnect. In the embodiment shown, various components can be coupled through an embedded controller **1435**. Such components can include a keyboard **1436** (e.g.,

coupled via a PS2 interface), a fan **1437**, and a thermal sensor **1439**. In some embodiments, touch pad **1430** may also couple to EC **1435** via a PS2 interface. In addition, a security processor such as a trusted platform module (TPM) **1438** may also couple to processor **1410** via this LPC interconnect.

(127) System **1400** can communicate with external devices in a variety of manners, including wirelessly. In the embodiment shown in FIG. **14**, various wireless modules, each of which can correspond to a radio configured for a particular wireless communication protocol, are present. One manner for wireless communication in a short range such as a near field may be via a NFC unit **1445** which may communicate, in one embodiment with processor **1410** via an SMBus. Note that via this NFC unit **1445**, devices in close proximity to each other can communicate.

(128) As further seen in FIG. **14**, additional wireless units can include other short range wireless engines including a WLAN unit **1450** and a Bluetooth™ unit **1452**. Using WLAN unit **1450**, Wi-Fi™ communications can be realized, while via Bluetooth™ unit **1452**, short range Bluetooth™ communications can occur. These units may communicate with processor **1410** via a given link.

(129) In addition, wireless wide area communications, e.g., according to a cellular or other wireless wide area protocol, can occur via a WWAN unit **1456** which in turn may couple to a subscriber identity module (SIM) **1457**. In addition, to enable receipt and use of location information, a GPS module **1455** may also be present. Note that in the embodiment shown in FIG. **14**, WWAN unit **1456** and an integrated capture device such as a camera module **1454** may communicate via a given link.

(130) To provide for audio inputs and outputs, an audio processor can be implemented via a digital signal processor (DSP) **1460**, which may couple to processor **1410** via a high definition audio (HDA) link. Similarly, DSP **1460** may communicate with an integrated coder/decoder (CODEC) and amplifier **1462** that in turn may couple to output speakers **1463** which may be implemented within the chassis. Similarly, amplifier and CODEC **1462** can be coupled to receive audio inputs from a microphone **1465** which in an embodiment can be implemented via dual array microphones (such as a digital microphone array) to provide for high quality audio inputs to enable voice-activated control of various operations within the system. Note also that audio outputs can be provided from amplifier/CODEC **1462** to a headphone jack **1464**. Although shown with these particular components in the embodiment of FIG. **14**, understand the scope of the present invention is not limited in this regard.

(131) Embodiments may be implemented in many different system types. Referring now to FIG. **15A**, shown is a block diagram of a system in accordance with an embodiment of the present invention. As shown in FIG. **15A**, multiprocessor system **1500** is a point-to-point interconnect system, and includes a first processor **1570** and a second processor **1580** coupled via a point-to-point interconnect **1550**. As shown in FIG. **15A**, each of processors **1570** and **1580** may be multicore processors, including first and second processor cores (i.e., processor cores **1574a** and **1574b** and processor cores **1584a** and **1584b**), although potentially many more cores may be present in the processors. Each of the processors can include a PCU or other power management logic to perform processor-based power management as described herein.

(132) Still referring to FIG. **15A**, first processor **1570** further includes an integrated memory controller (IMC) **1572** and point-to-point (P-P) interfaces **1576** and **1578**. Similarly, second processor **1580** includes an IMC **1582** and P-P interfaces **1586** and **1588**. As shown in FIG. **15**, IMCs **1572** and **1582** couple the processors to respective memories, namely a memory **1532** and a memory **1534**, which may be portions of system memory (e.g., DRAM) locally attached to the respective processors. First processor **1570** and second processor **1580** may be coupled to a chipset **1590** via P-P interconnects **1562** and **1564**, respectively. As shown in FIG. **15A**, chipset **1590** includes P-P interfaces **1594** and **1598**.

(133) Furthermore, chipset **1590** includes an interface **1592** to couple chipset **1590** with a high-performance graphics engine **1538**, by a P-P interconnect **1539**. In turn, chipset **1590** may be

coupled to a first bus **1516** via an interface **1596**. As shown in FIG. **15A**, various input/output (I/O) devices **1514** may be coupled to first bus **1516**, along with a bus bridge **1518** which couples first bus **1516** to a second bus **1520**. Various devices may be coupled to second bus **1520** including, for example, a keyboard/mouse **1522**, communication devices **1526** and a data storage unit **1528** such as a disk drive or other mass storage device which may include code **1530**, in one embodiment. Further, an audio I/O **1524** may be coupled to second bus **1520**. Embodiments can be incorporated into other types of systems including mobile devices such as a smart cellular telephone, tablet computer, netbook, Ultrabook™, or so forth.

(134) Referring now to FIG. **15B**, shown is a block diagram of a second more specific exemplary system **1501** in accordance with an embodiment of the present invention. Like elements in FIG. **15A** and FIG. **15B** bear like reference numerals, and certain aspects of FIG. **15A** have been omitted from FIG. **15B** in order to avoid obscuring other aspects of FIG. **15B**.

(135) FIG. **15B** illustrates that the processors **1570**, **1580** may include integrated memory and I/O control logic (“CL”) **1571** and **1581**, respectively. Thus, the control logic **1571** and **1581** include integrated memory controller units and include I/O control logic. FIG. **15N** illustrates that not only are the memories **1532**, **1534** coupled to the control logic **1571** and **1581**, but also that I/O devices **1513** are also coupled to the control logic **1571** and **1581**. Legacy I/O devices **1515** are coupled to the chipset **1590**.

(136) One or more aspects of at least one embodiment may be implemented by representative code stored on a machine-readable medium which represents and/or defines logic within an integrated circuit such as a processor. For example, the machine-readable medium may include instructions which represent various logic within the processor. When read by a machine, the instructions may cause the machine to fabricate the logic to perform the techniques described herein. Such representations, known as “IP cores,” are reusable units of logic for an integrated circuit that may be stored on a tangible, machine-readable medium as a hardware model that describes the structure of the integrated circuit. The hardware model may be supplied to various customers or manufacturing facilities, which load the hardware model on fabrication machines that manufacture the integrated circuit. The integrated circuit may be fabricated such that the circuit performs operations described in association with any of the embodiments described herein.

(137) FIG. **16** is a block diagram illustrating an IP core development system **1600** that may be used to manufacture an integrated circuit to perform operations according to an embodiment. The IP core development system **1600** may be used to generate modular, re-usable designs that can be incorporated into a larger design or used to construct an entire integrated circuit (e.g., an SoC integrated circuit). A design facility **1630** can generate a software simulation **1610** of an IP core design in a high-level programming language (e.g., C/C++). The software simulation **1610** can be used to design, test, and verify the behavior of the IP core. A register transfer level (RTL) design can then be created or synthesized from the simulation model. The RTL design **1615** is an abstraction of the behavior of the integrated circuit that models the flow of digital signals between hardware registers, including the associated logic performed using the modeled digital signals. In addition to an RTL design **1615**, lower-level designs at the logic level or transistor level may also be created, designed, or synthesized. Thus, the particular details of the initial design and simulation may vary.

(138) The RTL design **1615** or equivalent may be further synthesized by the design facility into a hardware model **1620**, which may be in a hardware description language (HDL), or some other representation of physical design data. The HDL may be further simulated or tested to verify the IP core design. The IP core design can be stored for delivery to a third-party fabrication facility **1665** using non-volatile memory **1640** (e.g., hard disk, flash memory, or any non-volatile storage medium). Alternately, the IP core design may be transmitted (e.g., via the Internet) over a wired connection **1650** or wireless connection **1660**. The fabrication facility **1665** may then fabricate an integrated circuit that is based at least in part on the IP core design. The fabricated integrated circuit

can be configured to perform operations in accordance with the components and/or processes described herein.

(139) FIGS. **17A-25** described below detail exemplary architectures and systems to implement embodiments of the components and/or processes described herein. In some embodiments, one or more hardware components and/or instructions described herein are emulated as detailed below, or are implemented as software modules.

(140) Embodiments of the instruction(s) detailed above are embodied may be embodied in a “generic vector friendly instruction format” which is detailed below. In other embodiments, such a format is not utilized and another instruction format is used, however, the description below of the writemask registers, various data transformations (swizzle, broadcast, etc.), addressing, etc. is generally applicable to the description of the embodiments of the instruction(s) above. Additionally, exemplary systems, architectures, and pipelines are detailed below. Embodiments of the instruction(s) above may be executed on such systems, architectures, and pipelines, but are not limited to those detailed.

(141) An instruction set may include one or more instruction formats. A given instruction format may define various fields (e.g., number of bits, location of bits) to specify, among other things, the operation to be performed (e.g., opcode) and the operand(s) on which that operation is to be performed and/or other data field(s) (e.g., mask). Some instruction formats are further broken down though the definition of instruction templates (or subformats). For example, the instruction templates of a given instruction format may be defined to have different subsets of the instruction format's fields (the included fields are typically in the same order, but at least some have different bit positions because there are less fields included) and/or defined to have a given field interpreted differently. Thus, each instruction of an ISA is expressed using a given instruction format (and, if defined, in a given one of the instruction templates of that instruction format) and includes fields for specifying the operation and the operands. For example, an exemplary ADD instruction has a specific opcode and an instruction format that includes an opcode field to specify that opcode and operand fields to select operands (source1/destination and source2); and an occurrence of this ADD instruction in an instruction stream will have specific contents in the operand fields that select specific operands. A set of SIMD extensions referred to as the Advanced Vector Extensions (AVX) (AVX1 and AVX2) and using the Vector Extensions (VEX) coding scheme has been released and/or published (e.g., see Intel® 64 and IA-32 Architectures Software Developer's Manual, September 2014; and see Intel® Advanced Vector Extensions Programming Reference, October 2014).

(142) Exemplary Instruction Formats

(143) Embodiments of the instruction(s) described herein may be embodied in different formats. Additionally, exemplary systems, architectures, and pipelines are detailed below. Embodiments of the instruction(s) may be executed on such systems, architectures, and pipelines, but are not limited to those detailed.

(144) Generic Vector Friendly Instruction Format

(145) A vector friendly instruction format is an instruction format that is suited for vector instructions (e.g., there are certain fields specific to vector operations). While embodiments are described in which both vector and scalar operations are supported through the vector friendly instruction format, alternative embodiments use only vector operations the vector friendly instruction format.

(146) FIGS. **17A-17B** are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to embodiments of the invention. FIG. **17A** is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to embodiments of the invention; while FIG. **17B** is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to embodiments of the invention. Specifically, a generic vector friendly instruction format **1700** for

which are defined class A and class B instruction templates, both of which include no memory access **1705** instruction templates and memory access **1720** instruction templates. The term generic in the context of the vector friendly instruction format refers to the instruction format not being tied to any specific instruction set.

(147) While embodiments of the invention will be described in which the vector friendly instruction format supports the following: a 64 byte vector operand length (or size) with 32 bit (4 byte) or 64 bit (8 byte) data element widths (or sizes) (and thus, a 64 byte vector consists of either 16 doubleword-size elements or alternatively, 8 quadword-size elements); a 64 byte vector operand length (or size) with 16 bit (2 byte) or 8 bit (1 byte) data element widths (or sizes); a 32 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); and a 16 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); alternative embodiments may support more, less and/or different vector operand sizes (e.g., 256 byte vector operands) with more, less, or different data element widths (e.g., 128 bit (16 byte) data element widths).

(148) The class A instruction templates in FIG. **17A** include: 1) within the no memory access **1705** instruction templates there is shown a no memory access, full round control type operation **1710** instruction template and a no memory access, data transform type operation **1715** instruction template; and 2) within the memory access **1720** instruction templates there is shown a memory access, temporal **1725** instruction template and a memory access, non-temporal **1730** instruction template. The class B instruction templates in FIG. **17B** include: 1) within the no memory access **1705** instruction templates there is shown a no memory access, write mask control, partial round control type operation **1712** instruction template and a no memory access, write mask control, vsize type operation **1717** instruction template; and 2) within the memory access **1720** instruction templates there is shown a memory access, write mask control **1727** instruction template.

(149) The generic vector friendly instruction format **1700** includes the following fields listed below in the order illustrated in FIGS. **17A-17B**.

(150) Format field **1740**—a specific value (an instruction format identifier value) in this field uniquely identifies the vector friendly instruction format, and thus occurrences of instructions in the vector friendly instruction format in instruction streams. As such, this field is optional in the sense that it is not needed for an instruction set that has only the generic vector friendly instruction format.

(151) Base operation field **1742**—its content distinguishes different base operations.

(152) Register index field **1744**—its content, directly or through address generation, specifies the locations of the source and destination operands, be they in registers or in memory. These include a sufficient number of bits to select N registers from a PxQ (e.g. 32×512, 16×128, 32×1024, 64×1024) register file. While in one embodiment N may be up to three sources and one destination register, alternative embodiments may support more or less sources and destination registers (e.g., may support up to two sources where one of these sources also acts as the destination, may support up to three sources where one of these sources also acts as the destination, may support up to two sources and one destination).

(153) Modifier field **1746**—its content distinguishes occurrences of instructions in the generic vector instruction format that specify memory access from those that do not; that is, between no memory access **1705** instruction templates and memory access **1720** instruction templates. Memory access operations read and/or write to the memory hierarchy (in some cases specifying the source and/or destination addresses using values in registers), while non-memory access operations do not (e.g., the source and destinations are registers). While in one embodiment this field also selects between three different ways to perform memory address calculations, alternative embodiments may support more, less, or different ways to perform memory address calculations.

(154) Augmentation operation field **1750**—its content distinguishes which one of a variety of different operations to be performed in addition to the base operation. This field is context specific.

In one embodiment of the invention, this field is divided into a class field **1768**, an alpha field **1752**, and a beta field **1754**. The augmentation operation field **1750** allows common groups of operations to be performed in a single instruction rather than 2, 3, or 4 instructions.

(155) Scale field **1760**—its content allows for the scaling of the index field's content for memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base}$).

(156) Displacement Field **1762A**—its content is used as part of memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{displacement}$).

(157) Displacement Factor Field **1762B** (note that the juxtaposition of displacement field **1762A** directly over displacement factor field **1762B** indicates one or the other is used)—its content is used as part of address generation; it specifies a displacement factor that is to be scaled by the size of a memory access (N)—where N is the number of bytes in the memory access (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{scaled displacement}$). Redundant low-order bits are ignored and hence, the displacement factor field's content is multiplied by the memory operands total size (N) in order to generate the final displacement to be used in calculating an effective address. The value of N is determined by the processor hardware at runtime based on the full opcode field **1774** (described later herein) and the data manipulation field **1754C**. The displacement field **1762A** and the displacement factor field **1762B** are optional in the sense that they are not used for the no memory access **1705** instruction templates and/or different embodiments may implement only one or none of the two.

(158) Data element width field **1764**—its content distinguishes which one of a number of data element widths is to be used (in some embodiments for all instructions; in other embodiments for only some of the instructions). This field is optional in the sense that it is not needed if only one data element width is supported and/or data element widths are supported using some aspect of the opcodes.

(159) Write mask field **1770**—its content controls, on a per data element position basis, whether that data element position in the destination vector operand reflects the result of the base operation and augmentation operation. Class A instruction templates support merging-writemasking, while class B instruction templates support both merging- and zeroing-writemasking. When merging, vector masks allow any set of elements in the destination to be protected from updates during the execution of any operation (specified by the base operation and the augmentation operation); in other one embodiment, preserving the old value of each element of the destination where the corresponding mask bit has a 0. In contrast, when zeroing vector masks allow any set of elements in the destination to be zeroed during the execution of any operation (specified by the base operation and the augmentation operation); in one embodiment, an element of the destination is set to 0 when the corresponding mask bit has a 0 value. A subset of this functionality is the ability to control the vector length of the operation being performed (that is, the span of elements being modified, from the first to the last one); however, it is not necessary that the elements that are modified be consecutive. Thus, the write mask field **1770** allows for partial vector operations, including loads, stores, arithmetic, logical, etc. While embodiments of the invention are described in which the write mask field's **1770** content selects one of a number of write mask registers that contains the write mask to be used (and thus the write mask field's **1770** content indirectly identifies that masking to be performed), alternative embodiments instead or additionally allow the mask write field's **1770** content to directly specify the masking to be performed.

(160) Immediate field **1772**—its content allows for the specification of an immediate. This field is optional in the sense that it is not present in an implementation of the generic vector friendly format that does not support immediate and it is not present in instructions that do not use an immediate.

(161) Class field **1768**—its content distinguishes between different classes of instructions. With reference to FIGS. **17A-B**, the contents of this field select between class A and class B instructions. In FIGS. **17A-B**, rounded corner squares are used to indicate a specific value is present in a field (e.g., class A **1768A** and class B **1768B** for the class field **1768** respectively in FIGS. **17A-B**).

(162) Instruction Templates of Class A

(163) In the case of the non-memory access **1705** instruction templates of class A, the alpha field **1752** is interpreted as an RS field **1752A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **1752A.1** and data transform **1752A.2** are respectively specified for the no memory access, round type operation **1710** and the no memory access, data transform type operation **1715** instruction templates), while the beta field **1754** distinguishes which of the operations of the specified type is to be performed. In the no memory access **1705** instruction templates, the scale field **1760**, the displacement field **1762A**, and the displacement scale field **1762B** are not present.

(164) No-Memory Access Instruction Templates—Full Round Control Type Operation

(165) In the no memory access full round control type operation **1710** instruction template, the beta field **1754** is interpreted as a round control field **1754A**, whose content(s) provide static rounding. While in the described embodiments of the invention the round control field **1754A** includes a suppress all floating point exceptions (SAE) field **1756** and a round operation control field **1758**, alternative embodiments may support may encode both these concepts into the same field or only have one or the other of these concepts/fields (e.g., may have only the round operation control field **1758**).

(166) SAE field **1756**—its content distinguishes whether or not to disable the exception event reporting; when the SAE field's **1756** content indicates suppression is enabled, a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler.

(167) Round operation control field **1758**—its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **1758** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the invention where a processor includes a control register for specifying rounding modes, the round operation control field's **1750** content overrides that register value.

(168) No Memory Access Instruction Templates—Data Transform Type Operation

(169) In the no memory access data transform type operation **1715** instruction template, the beta field **1754** is interpreted as a data transform field **1754B**, whose content distinguishes which one of a number of data transforms is to be performed (e.g., no data transform, swizzle, broadcast).

(170) In the case of a memory access **1720** instruction template of class A, the alpha field **1752** is interpreted as an eviction hint field **1752B**, whose content distinguishes which one of the eviction hints is to be used (in FIG. 17A, temporal **1752B.1** and non-temporal **1752B.2** are respectively specified for the memory access, temporal **1725** instruction template and the memory access, non-temporal **1730** instruction template), while the beta field **1754** is interpreted as a data manipulation field **1754C**, whose content distinguishes which one of a number of data manipulation operations (also known as primitives) is to be performed (e.g., no manipulation; broadcast; up conversion of a source; and down conversion of a destination). The memory access **1720** instruction templates include the scale field **1760**, and optionally the displacement field **1762A** or the displacement scale field **1762B**.

(171) Vector memory instructions perform vector loads from and vector stores to memory, with conversion support. As with regular vector instructions, vector memory instructions transfer data from/to memory in a data element-wise fashion, with the elements that are actually transferred is dictated by the contents of the vector mask that is selected as the write mask.

(172) Memory Access Instruction Templates—Temporal

(173) Temporal data is data likely to be reused soon enough to benefit from caching. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

(174) Memory Access Instruction Templates—Non-Temporal

(175) Non-temporal data is data unlikely to be reused soon enough to benefit from caching in the 1st-level cache and should be given priority for eviction. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

(176) Instruction Templates of Class B

(177) In the case of the instruction templates of class B, the alpha field **1752** is interpreted as a write mask control (Z) field **1752C**, whose content distinguishes whether the write masking controlled by the write mask field **1770** should be a merging or a zeroing.

(178) In the case of the non-memory access **1705** instruction templates of class B, part of the beta field **1754** is interpreted as an RL field **1757A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **1757A.1** and vector length (VSIZE) **1757A.2** are respectively specified for the no memory access, write mask control, partial round control type operation **1712** instruction template and the no memory access, write mask control, VSIZE type operation **1717** instruction template), while the rest of the beta field **1754** distinguishes which of the operations of the specified type is to be performed. In the no memory access **1705** instruction templates, the scale field **1760**, the displacement field **1762A**, and the displacement scale field **1762B** are not present.

(179) In the no memory access, write mask control, partial round control type operation **1710** instruction template, the rest of the beta field **1754** is interpreted as a round operation field **1759A** and exception event reporting is disabled (a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler).

(180) Round operation control field **1759A**—just as round operation control field **1758**, its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **1759A** allows for the changing of the rounding mode on a per instruction basis. In one embodiment of the invention where a processor includes a control register for specifying rounding modes, the round operation control field's **1750** content overrides that register value.

(181) In the no memory access, write mask control, VSIZE type operation **1717** instruction template, the rest of the beta field **1754** is interpreted as a vector length field **1759B**, whose content distinguishes which one of a number of data vector lengths is to be performed on (e.g., 128, 256, or 512 byte).

(182) In the case of a memory access **1720** instruction template of class B, part of the beta field **1754** is interpreted as a broadcast field **1757B**, whose content distinguishes whether or not the broadcast type data manipulation operation is to be performed, while the rest of the beta field **1754** is interpreted the vector length field **1759B**. The memory access **1720** instruction templates include the scale field **1760**, and optionally the displacement field **1762A** or the displacement scale field **1762B**.

(183) With regard to the generic vector friendly instruction format **1700**, a full opcode field **1774** is shown including the format field **1740**, the base operation field **1742**, and the data element width field **1764**. While one embodiment is shown where the full opcode field **1774** includes all of these fields, the full opcode field **1774** includes less than all of these fields in embodiments that do not support all of them. The full opcode field **1774** provides the operation code (opcode).

(184) The augmentation operation field **1750**, the data element width field **1764**, and the write mask field **1770** allow these features to be specified on a per instruction basis in the generic vector friendly instruction format.

(185) The combination of write mask field and data element width field create typed instructions in that they allow the mask to be applied based on different data element widths.

(186) The various instruction templates found within class A and class B are beneficial in different situations. In some embodiments of the invention, different processors or different cores within a processor may support only class A, only class B, or both classes. For instance, a high performance general purpose out-of-order core intended for general-purpose computing may support only class

B, a core intended primarily for graphics and/or scientific (throughput) computing may support only class A, and a core intended for both may support both (of course, a core that has some mix of templates and instructions from both classes but not all templates and instructions from both classes is within the purview of the invention). Also, a single processor may include multiple cores, all of which support the same class or in which different cores support different class. For instance, in a processor with separate graphics and general purpose cores, one of the graphics cores intended primarily for graphics and/or scientific computing may support only class A, while one or more of the general purpose cores may be high performance general purpose cores with out of order execution and register renaming intended for general-purpose computing that support only class B. Another processor that does not have a separate graphics core, may include one more general purpose in-order or out-of-order cores that support both class A and class B. Of course, features from one class may also be implement in the other class in different embodiments of the invention. Programs written in a high level language would be put (e.g., just in time compiled or statically compiled) into an variety of different executable forms, including: 1) a form having only instructions of the class(es) supported by the target processor for execution; or 2) a form having alternative routines written using different combinations of the instructions of all classes and having control flow code that selects the routines to execute based on the instructions supported by the processor which is currently executing the code.

(187) Exemplary Specific Vector Friendly Instruction Format

(188) FIG. **18A-18C** are block diagrams illustrating an exemplary specific vector friendly instruction format according to embodiments of the invention. FIG. **18A** shows a specific vector friendly instruction format **1800** that is specific in the sense that it specifies the location, size, interpretation, and order of the fields, as well as values for some of those fields. The specific vector friendly instruction format **1800** may be used to extend the x86 instruction set, and thus some of the fields are similar or the same as those used in the existing x86 instruction set and extension thereof (e.g., AVX). This format remains consistent with the prefix encoding field, real opcode byte field, MOD R/M field, SIB field, displacement field, and immediate fields of the existing x86 instruction set with extensions. The fields from FIGS. **17A-17B** into which the fields from FIGS. **18A-18C** map are illustrated.

(189) It should be understood that, although embodiments of the invention are described with reference to the specific vector friendly instruction format **1800** in the context of the generic vector friendly instruction format **1700** for illustrative purposes, the invention is not limited to the specific vector friendly instruction format **1800** except where claimed. For example, the generic vector friendly instruction format **1700** contemplates a variety of possible sizes for the various fields, while the specific vector friendly instruction format **1800** is shown as having fields of specific sizes. By way of specific example, while the data element width field **1764** is illustrated as a one bit field in the specific vector friendly instruction format **1800**, the invention is not so limited (that is, the generic vector friendly instruction format **1700** contemplates other sizes of the data element width field **1764**).

(190) The generic vector friendly instruction format **1700** includes the following fields listed below in the order illustrated in FIG. **18A**.

(191) EVEX Prefix (Bytes **0-3**) **1802**—is encoded in a four-byte form.

(192) Format Field **1740** (EVEX Byte **0**, bits [7:0])—the first byte (EVEX Byte **0**) is the format field **1740** and it contains 0x62 (the unique value used for distinguishing the vector friendly instruction format in one embodiment of the invention).

(193) The second-fourth bytes (EVEX Bytes **1-3**) include a number of bit fields providing specific capability.

(194) REX field **1805** (EVEX Byte **1**, bits [7-5])—consists of a EVEX.R bit field (EVEX Byte **1**, bit [7]—R), EVEX.X bit field (EVEX byte **1**, bit [6]—X), and EVEX.B byte **1**, bit[5]—B). The EVEX.R, EVEX.X, and EVEX.B bit fields provide the same functionality as the corresponding

VEX bit fields, and are encoded using 1 s complement form, i.e. ZMM0 is encoded as 1111B, ZMM15 is encoded as 0000B. Other fields of the instructions encode the lower three bits of the register indexes as is known in the art (rrr, xxx, and bbb), so that Rrrr, Xxxx, and Bbbb may be formed by adding EVEX.R, EVEX.X, and EVEX.B.

(195) REX' field **1810**—this is the first part of the REX' field **1810** and is the EVEX.R' bit field (EVEX Byte **1**, bit [4]—R') that is used to encode either the upper 16 or lower 16 of the extended 32 register set. In one embodiment of the invention, this bit, along with others as indicated below, is stored in bit inverted format to distinguish (in the well-known x86 32-bit mode) from the BOUND instruction, whose real opcode byte is 62, but does not accept in the MOD R/M field (described below) the value of 11 in the MOD field; alternative embodiments of the invention do not store this and the other indicated bits below in the inverted format. A value of 1 is used to encode the lower 16 registers. In other words, R'Rrrr is formed by combining EVEX.R', EVEX.R, and the other RRR from other fields.

(196) Opcode map field **1815** (EVEX byte **1**, bits [3:0]—mmmm)—its content encodes an implied leading opcode byte (**0F**, **0F 38**, or **0F 3**).

(197) Data element width field **1764** (EVEX byte **2**, bit [7]—W)—is represented by the notation EVEX.W. EVEX.W is used to define the granularity (size) of the datatype (either 32-bit data elements or 64-bit data elements).

(198) EVEX.vvvv **1820** (EVEX Byte **2**, bits [6:3]—vvvv)—the role of EVEX.vvvv may include the following: 1) EVEX.vvvv encodes the first source register operand, specified in inverted (1 s complement) form and is valid for instructions with 2 or more source operands; 2) EVEX.vvvv encodes the destination register operand, specified in 1 s complement form for certain vector shifts; or 3) EVEX.vvvv does not encode any operand, the field is reserved and should contain **1111b**. Thus, EVEX.vvvv field **1820** encodes the 4 low-order bits of the first source register specifier stored in inverted (1 s complement) form. Depending on the instruction, an extra different EVEX bit field is used to extend the specifier size to 32 registers.

(199) EVEX. U **1768** Class field (EVEX byte **2**, bit [2]—U)—If EVEX.U=0, it indicates class A or EVEX.U0; if EVEX.U=1, it indicates class B or EVEX.U1.

(200) Prefix encoding field **1825** (EVEX byte **2**, bits [1:0]—pp)—provides additional bits for the base operation field. In addition to providing support for the legacy SSE instructions in the EVEX prefix format, this also has the benefit of compacting the SIMD prefix (rather than requiring a byte to express the SIMD prefix, the EVEX prefix requires only 2 bits). In one embodiment, to support legacy SSE instructions that use a SIMD prefix (66H, F2H, F3H) in both the legacy format and in the EVEX prefix format, these legacy SIMD prefixes are encoded into the SIMD prefix encoding field; and at runtime are expanded into the legacy SIMD prefix prior to being provided to the decoder's PLA (so the PLA can execute both the legacy and EVEX format of these legacy instructions without modification). Although newer instructions could use the EVEX prefix encoding field's content directly as an opcode extension, certain embodiments expand in a similar fashion for consistency but allow for different meanings to be specified by these legacy SIMD prefixes. An alternative embodiment may redesign the PLA to support the 2 bit SIMD prefix encodings, and thus not require the expansion.

(201) Alpha field **1752** (EVEX byte **3**, bit [7]—EH; also known as EVEX.EH, EVEX.rs, EVEX.RL, EVEX.write mask control, and EVEX.N; also illustrated with a)—as previously described, this field is context specific.

(202) Beta field **1754** (EVEX byte **3**, bits [6:4]—SSS, also known as EVEX.s.sub.2-0, EVEX.r.sub.2-0, EVEX.rr1, EVEX.LL0, EVEX.LLB; also illustrated with $\beta\beta\beta$)—as previously described, this field is context specific.

(203) REX' field **1810**—this is the remainder of the REX' field and is the EVEX.V' bit field (EVEX Byte **3**, bit [3]—V') that may be used to encode either the upper 16 or lower 16 of the extended 32 register set. This bit is stored in bit inverted format. A value of 1 is used to encode the

lower 16 registers. In other words, V' VVVV is formed by combining EVEX.V', EVEX.vvvv.

(204) Write mask field **1770** (EVEX byte 3, bits [2:0]-kkk)—its content specifies the index of a register in the write mask registers as previously described. In one embodiment of the invention, the specific value EVEX.kkk=000 has a special behavior implying no write mask is used for the particular instruction (this may be implemented in a variety of ways including the use of a write mask hardwired to all ones or hardware that bypasses the masking hardware).

(205) Real Opcode Field **1830** (Byte 4) is also known as the opcode byte. Part of the opcode is specified in this field.

(206) MOD R/M Field **1840** (Byte 5) includes MOD field **1842**, Reg field **1844**, and R/M field **1846**. As previously described, the MOD field's **1842** content distinguishes between memory access and non-memory access operations. The role of Reg field **1844** can be summarized to two situations: encoding either the destination register operand or a source register operand, or be treated as an opcode extension and not used to encode any instruction operand. The role of R/M field **1846** may include the following: encoding the instruction operand that references a memory address, or encoding either the destination register operand or a source register operand.

(207) Scale, Index, Base (SIB) Byte (Byte 6)—As previously described, the scale field's **1850** content is used for memory address generation. SIB.xxx **1854** and SIB.bbb **1856**—the contents of these fields have been previously referred to with regard to the register indexes Xxxx and Bbbb.

(208) Displacement field **1762A** (Bytes 7-10)—when MOD field **1842** contains 10, bytes 7-10 are the displacement field **1762A**, and it works the same as the legacy 32-bit displacement (disp32) and works at byte granularity.

(209) Displacement factor field **1762B** (Byte 7)—when MOD field **1842** contains 01, byte 7 is the displacement factor field **1762B**. The location of this field is that same as that of the legacy x86 instruction set 8-bit displacement (disp8), which works at byte granularity. Since disp8 is sign extended, it can only address between -128 and 127 bytes offsets; in terms of 64 byte cache lines, disp8 uses 8 bits that can be set to only four really useful values -128, -64, 0, and 64; since a greater range is often needed, disp32 is used; however, disp32 requires 4 bytes. In contrast to disp8 and disp32, the displacement factor field **1762B** is a reinterpretation of disp8; when using displacement factor field **1762B**, the actual displacement is determined by the content of the displacement factor field multiplied by the size of the memory operand access (N). This type of displacement is referred to as disp8*N. This reduces the average instruction length (a single byte of used for the displacement but with a much greater range). Such compressed displacement is based on the assumption that the effective displacement is multiple of the granularity of the memory access, and hence, the redundant low-order bits of the address offset do not need to be encoded. In other words, the displacement factor field **1762B** substitutes the legacy x86 instruction set 8-bit displacement. Thus, the displacement factor field **1762B** is encoded the same way as an x86 instruction set 8-bit displacement (so no changes in the ModRM/SIB encoding rules) with the only exception that disp8 is overloaded to disp8*N. In other words, there are no changes in the encoding rules or encoding lengths but only in the interpretation of the displacement value by hardware (which needs to scale the displacement by the size of the memory operand to obtain a byte-wise address offset). Immediate field **1772** operates as previously described.

(210) Full Opcode Field

(211) FIG. **18B** is a block diagram illustrating the fields of the specific vector friendly instruction format **1800** that make up the full opcode field **1774** according to one embodiment of the invention. Specifically, the full opcode field **1774** includes the format field **1740**, the base operation field **1742**, and the data element width (W) field **1764**. The base operation field **1742** includes the prefix encoding field **1825**, the opcode map field **1815**, and the real opcode field **1830**.

(212) Register Index Field

(213) FIG. **18C** is a block diagram illustrating the fields of the specific vector friendly instruction format **1800** that make up the register index field **1744** according to one embodiment of the

invention. Specifically, the register index field **1744** includes the REX field **1805**, the REX' field **1810**, the MODR/M.reg field **1844**, the MODR/M.r/m field **1846**, the VVVV field **1820**, xxx field **1854**, and the bbb field **1856**.

(214) Augmentation Operation Field

(215) FIG. **18D** is a block diagram illustrating the fields of the specific vector friendly instruction format **1800** that make up the augmentation operation field **1750** according to one embodiment of the invention. When the class (U) field **1768** contains 0, it signifies EVEX.U0 (class A **1768A**); when it contains 1, it signifies EVEX.U1 (class B **1768B**). When U=0 and the MOD field **1842** contains 11 (signifying a no memory access operation), the alpha field **1752** (EVEX byte 3, bit [7]—EH) is interpreted as the rs field **1752A**. When the rs field **1752A** contains a 1 (round **1752A.1**), the beta field **1754** (EVEX byte 3, bits [6:4]—SSS) is interpreted as the round control field **1754A**. The round control field **1754A** includes a one bit SAE field **1756** and a two bit round operation field **1758**. When the rs field **1752A** contains a 0 (data transform **1752A.2**), the beta field **1754** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data transform field **1754B**. When U=0 and the MOD field **1842** contains 00, 01, or 10 (signifying a memory access operation), the alpha field **1752** (EVEX byte 3, bit [7]—EH) is interpreted as the eviction hint (EH) field **1752B** and the beta field **1754** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data manipulation field **1754C**.

(216) When U=1, the alpha field **1752** (EVEX byte 3, bit [7]—EH) is interpreted as the write mask control (Z) field **1752C**. When U=1 and the MOD field **1842** contains 11 (signifying a no memory access operation), part of the beta field **1754** (EVEX byte 3, bit [4]—S.sub.0) is interpreted as the RL field **1757A**; when it contains a 1 (round **1757A.1**) the rest of the beta field **1754** (EVEX byte 3, bit [6-5]-S.sub.2-1) is interpreted as the round operation field **1759A**, while when the RL field **1757A** contains a 0 (VSIZE **1757.A2**) the rest of the beta field **1754** (EVEX byte 3, bit [6-5]-S.sub.2-1) is interpreted as the vector length field **1759B** (EVEX byte 3, bit [6-5]-L.sub.1-0). When U=1 and the MOD field **1842** contains 00, 01, or 10 (signifying a memory access operation), the beta field **1754** (EVEX byte 3, bits [6:4]-SSS) is interpreted as the vector length field **1759B** (EVEX byte 3, bit [6-5]—L.sub.1-0) and the broadcast field **1757B** (EVEX byte 3, bit [4]—B).

(217) Exemplary Register Architecture

(218) FIG. **19** is a block diagram of a register architecture **1900** according to one embodiment of the invention. In the embodiment illustrated, there are 32 vector registers **1910** that are 512 bits wide; these registers are referenced as zmm0 through zmm31. The lower order 256 bits of the lower 16 zmm registers are overlaid on registers ymm0-16. The lower order 128 bits of the lower 16 zmm registers (the lower order 128 bits of the ymm registers) are overlaid on registers xmm0-15. The specific vector friendly instruction format **1800** operates on these overlaid register file as illustrated in the below tables.

(219) TABLE-US-00001 Adjustable Vector Length Class Operations Registers Instruction A 1710, 1715, zmm registers (the Templates that (FIG. 17A; 1725, 1730 vector length is 64 do not include U = 0) byte) the vector B 1712 zmm registers (the length field (FIG. 17B; vector length is 64 1759B U = 1) byte) Instruction B 1717, 1727 zmm, ymm, or xmm templates that (FIG. 17B; registers (the vector do include U = 1) length is 64 byte, the vector 32 byte, or 16 byte) length field depending on the vector 1759B length field 1759B

(220) In other words, the vector length field **1759B** selects between a maximum length and one or more other shorter lengths, where each such shorter length is half the length of the preceding length; and instructions templates without the vector length field **1759B** operate on the maximum vector length. Further, in one embodiment, the class B instruction templates of the specific vector friendly instruction format **1800** operate on packed or scalar single/double-precision floating point data and packed or scalar integer data. Scalar operations are operations performed on the lowest order data element position in an zmm/ymm/xmm register; the higher order data element positions are either left the same as they were prior to the instruction or zeroed depending on the

embodiment.

(221) Write mask registers **1915**—in the embodiment illustrated, there are 8 write mask registers (**k0** through **k7**), each 64 bits in size. In an alternate embodiment, the write mask registers **1915** are 16 bits in size. As previously described, in one embodiment of the invention, the vector mask register **k0** cannot be used as a write mask; when the encoding that would normally indicate **k0** is used for a write mask, it selects a hardwired write mask of **0xFFFF**, effectively disabling write masking for that instruction.

(222) General-purpose registers **1925**—in the embodiment illustrated, there are sixteen 64-bit general-purpose registers that are used along with the existing x86 addressing modes to address memory operands. These registers are referenced by the names **RAX**, **RBX**, **RCX**, **RDX**, **RBP**, **RSI**, **RDI**, **RSP**, and **R8** through **R15**.

(223) Scalar floating point stack register file (x87 stack) **1945**, on which is aliased the MMX packed integer flat register file **1950**—in the embodiment illustrated, the x87 stack is an eight-element stack used to perform scalar floating-point operations on 32/64/80-bit floating point data using the x87 instruction set extension; while the MMX registers are used to perform operations on 64-bit packed integer data, as well as to hold operands for some operations performed between the MMX and XMM registers.

(224) Alternative embodiments of the invention may use wider or narrower registers. Additionally, alternative embodiments of the invention may use more, less, or different register files and registers.

(225) Exemplary Core Architectures, Processors, and Computer Architectures

(226) Processor cores may be implemented in different ways, for different purposes, and in different processors. For instance, implementations of such cores may include: 1) a general purpose in-order core intended for general-purpose computing; 2) a high performance general purpose out-of-order core intended for general-purpose computing; 3) a special purpose core intended primarily for graphics and/or scientific (throughput) computing. Implementations of different processors may include: 1) a CPU including one or more general purpose in-order cores intended for general-purpose computing and/or one or more general purpose out-of-order cores intended for general-purpose computing; and 2) a coprocessor including one or more special purpose cores intended primarily for graphics and/or scientific (throughput). Such different processors lead to different computer system architectures, which may include: 1) the coprocessor on a separate chip from the CPU; 2) the coprocessor on a separate die in the same package as a CPU; 3) the coprocessor on the same die as a CPU (in which case, such a coprocessor is sometimes referred to as special purpose logic, such as integrated graphics and/or scientific (throughput) logic, or as special purpose cores); and 4) a system on a chip that may include on the same die the described CPU (sometimes referred to as the application core(s) or application processor(s)), the above described coprocessor, and additional functionality. Exemplary core architectures are described next, followed by descriptions of exemplary processors and computer architectures.

(227) Exemplary Core Architectures

(228) In-Order and Out-of-Order Core Block Diagram

(229) FIG. **20A** is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to embodiments of the invention. FIG. **20B** is a block diagram illustrating both an exemplary embodiment of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to embodiments of the invention. The solid lined boxes in FIGS. **20A-B** illustrate the in-order pipeline and in-order core, while the optional addition of the dashed lined boxes illustrates the register renaming, out-of-order issue/execution pipeline and core. Given that the in-order aspect is a subset of the out-of-order aspect, the out-of-order aspect will be described.

(230) In FIG. **20A**, a processor pipeline **2000** includes a fetch stage **2002**, a length decode stage

2004, a decode stage **2006**, an allocation stage **2008**, a renaming stage **2010**, a scheduling (also known as a dispatch or issue) stage **2012**, a register read/memory read stage **2014**, an execute stage **2016**, a write back/memory write stage **2018**, an exception handling stage **2022**, and a commit stage **2024**.

(231) FIG. **20B** shows processor core **2090** including a front end unit **2030** coupled to an execution engine unit **2050**, and both are coupled to a memory unit **2070**. The core **2090** may be a reduced instruction set computing (RISC) core, a complex instruction set computing (CISC) core, a very long instruction word (VLIW) core, or a hybrid or alternative core type. As yet another option, the core **2090** may be a special-purpose core, such as, for example, a network or communication core, compression engine, coprocessor core, general purpose computing graphics processing unit (GPGPU) core, graphics core, or the like.

(232) The front end unit **2030** includes a branch prediction unit **2032** coupled to an instruction cache unit **2034**, which is coupled to an instruction translation lookaside buffer (TLB) **2036**, which is coupled to an instruction fetch unit **2038**, which is coupled to a decode unit **2040**. The decode unit **2040** (or decoder) may decode instructions, and generate as an output one or more micro-operations, micro-code entry points, microinstructions, other instructions, or other control signals, which are decoded from, or which otherwise reflect, or are derived from, the original instructions. The decode unit **2040** may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. In one embodiment, the core **2090** includes a microcode ROM or other medium that stores microcode for certain macroinstructions (e.g., in decode unit **2040** or otherwise within the front end unit **2030**). The decode unit **2040** is coupled to a rename/allocator unit **2052** in the execution engine unit **2050**.

(233) The execution engine unit **2050** includes the rename/allocator unit **2052** coupled to a retirement unit **2054** and a set of one or more scheduler unit(s) **2056**. The scheduler unit(s) **2056** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) **2056** is coupled to the physical register file(s) unit(s) **2058**. Each of the physical register file(s) units **2058** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point status (e.g., an instruction pointer that is the address of the next instruction to be executed), etc. In one embodiment, the physical register file(s) unit **2058** comprises a vector registers unit, a write mask registers unit, and a scalar registers unit. These register units may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) unit(s) **2058** is overlapped by the retirement unit **2054** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement unit **2054** and the physical register file(s) unit(s) **2058** are coupled to the execution cluster(s) **2060**. The execution cluster(s) **2060** includes a set of one or more execution units **2062** and a set of one or more memory access units **2064**. The execution units **2062** may perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some embodiments may include a number of execution units dedicated to specific functions or sets of functions, other embodiments may include only one execution unit or multiple execution units that all perform all functions. The scheduler unit(s) **2056**, physical register file(s) unit(s) **2058**, and execution cluster(s) **2060** are shown as being possibly plural because certain embodiments create separate pipelines for certain types of data/operations (e.g., a scalar integer pipeline, a scalar floating point/packed integer/packed floating point/vector integer/vector floating point pipeline, and/or a memory access pipeline that each have their own scheduler unit, physical register file(s) unit, and/or execution

cluster—and in the case of a separate memory access pipeline, certain embodiments are implemented in which only the execution cluster of this pipeline has the memory access unit(s) **2064**). It should also be understood that where separate pipelines are used, one or more of these pipelines may be out-of-order issue/execution and the rest in-order.

(234) The set of memory access units **2064** is coupled to the memory unit **2070**, which includes a data TLB unit **2072** coupled to a data cache unit **2074** coupled to a level 2 (L2) cache unit **2076**. In one exemplary embodiment, the memory access units **2064** may include a load unit, a store address unit, and a store data unit, each of which is coupled to the data TLB unit **2072** in the memory unit **2070**. The instruction cache unit **2034** is further coupled to a level 2 (L2) cache unit **2076** in the memory unit **2070**. The L2 cache unit **2076** is coupled to one or more other levels of cache and eventually to a main memory.

(235) By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement the pipeline **2000** as follows: 1) the instruction fetch **2038** performs the fetch and length decoding stages **2002** and **2004**; 2) the decode unit **2040** performs the decode stage **2006**; 3) the rename/allocator unit **2052** performs the allocation stage **2008** and renaming stage **2010**; 4) the scheduler unit(s) **2056** performs the schedule stage **2012**; 5) the physical register file(s) unit(s) **2058** and the memory unit **2070** perform the register read/memory read stage **2014**; the execution cluster **2060** perform the execute stage **2016**; 6) the memory unit **2070** and the physical register file(s) unit(s) **2058** perform the write back/memory write stage **2018**; 7) various units may be involved in the exception handling stage **2022**; and 8) the retirement unit **2054** and the physical register file(s) unit(s) **2058** perform the commit stage **2024**.

(236) The core **2090** may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one embodiment, the core **2090** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

(237) It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyperthreading technology).

(238) While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated embodiment of the processor also includes separate instruction and data cache units **2034/2074** and a shared L2 cache unit **2076**, alternative embodiments may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some embodiments, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

(239) Specific Exemplary In-Order Core Architecture

(240) FIGS. **21A-B** illustrate a block diagram of a more specific exemplary in-order core architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip. The logic blocks communicate through a high-bandwidth interconnect network (e.g., a ring network) with some fixed function logic, memory I/O interfaces, and other necessary I/O logic, depending on the application.

(241) FIG. **21A** is a block diagram of a single processor core, along with its connection to the on-die interconnect network **2102** and with its local subset of the Level 2 (L2) cache **2104**, according

to embodiments of the invention. In one embodiment, an instruction decoder **2100** supports the x86 instruction set with a packed data instruction set extension. An L1 cache **2106** allows low-latency accesses to cache memory into the scalar and vector units. While in one embodiment (to simplify the design), a scalar unit **2108** and a vector unit **2110** use separate register sets (respectively, scalar registers **2112** and vector registers **2114**) and data transferred between them is written to memory and then read back in from a level 1 (L1) cache **2106**, alternative embodiments of the invention may use a different approach (e.g., use a single register set or include a communication path that allow data to be transferred between the two register files without being written and read back). (242) The local subset of the L2 cache **2104** is part of a global L2 cache that is divided into separate local subsets, one per processor core. Each processor core has a direct access path to its own local subset of the L2 cache **2104**. Data read by a processor core is stored in its L2 cache subset **2104** and can be accessed quickly, in parallel with other processor cores accessing their own local L2 cache subsets. Data written by a processor core is stored in its own L2 cache subset **2104** and is flushed from other subsets, if necessary. The ring network ensures coherency for shared data. The ring network is bi-directional to allow agents such as processor cores, L2 caches and other logic blocks to communicate with each other within the chip. Each ring data-path is 1012-bits wide per direction.

(243) FIG. **21B** is an expanded view of part of the processor core in FIG. **21A** according to embodiments of the invention. FIG. **21B** includes an L1 data cache **2106A** part of the L1 cache **2104**, as well as more detail regarding the vector unit **2110** and the vector registers **2114**. Specifically, the vector unit **2110** is a 16-wide vector processing unit (VPU) (see the 16-wide ALU **2128**), which executes one or more of integer, single-precision float, and double-precision float instructions. The VPU supports swizzling the register inputs with swizzle unit **2120**, numeric conversion with numeric convert units **2122A-B**, and replication with replication unit **2124** on the memory input. Write mask registers **2126** allow predicating resulting vector writes.

(244) FIG. **22** is a block diagram of a processor **2200** that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to embodiments of the invention. The solid lined boxes in FIG. **22** illustrate a processor **2200** with a single core **2202A**, a system agent **2210**, a set of one or more bus controller units **2216**, while the optional addition of the dashed lined boxes illustrates an alternative processor **2200** with multiple cores **2202A-N**, a set of one or more integrated memory controller unit(s) **2214** in the system agent unit **2210**, and special purpose logic **2208**.

(245) Thus, different implementations of the processor **2200** may include: 1) a CPU with the special purpose logic **2208** being integrated graphics and/or scientific (throughput) logic (which may include one or more cores), and the cores **2202A-N** being one or more general purpose cores (e.g., general purpose in-order cores, general purpose out-of-order cores, a combination of the two); 2) a coprocessor with the cores **2202A-N** being a large number of special purpose cores intended primarily for graphics and/or scientific (throughput); and 3) a coprocessor with the cores **2202A-N** being a large number of general purpose in-order cores. Thus, the processor **2200** may be a general-purpose processor, coprocessor or special-purpose processor, such as, for example, a network or communication processor, compression engine, graphics processor, GPGPU (general purpose graphics processing unit), a high-throughput many integrated core (MIC) coprocessor (including 30 or more cores), embedded processor, or the like. The processor may be implemented on one or more chips. The processor **2200** may be a part of and/or may be implemented on one or more substrates using any of a number of process technologies, such as, for example, BiCMOS, CMOS, or NMOS.

(246) The memory hierarchy includes one or more levels of cache within the cores, a set or one or more shared cache units **2206**, and external memory (not shown) coupled to the set of integrated memory controller units **2214**. The set of shared cache units **2206** may include one or more mid-level caches, such as level 2 (L2), level 3 (L3), level 4 (L4), or other levels of cache, a last level

cache (LLC), and/or combinations thereof.

(247) While in one embodiment a ring based interconnect unit **2212** interconnects the integrated graphics logic **2208**, the set of shared cache units **2206**, and the system agent unit **2210**/integrated memory controller unit(s) **2214**, alternative embodiments may use any number of well-known techniques for interconnecting such units. In one embodiment, coherency is maintained between one or more cache units **2206** and cores **2202-A-N**.

(248) In some embodiments, one or more of the cores **2202A-N** are capable of multi-threading. The system agent **2210** includes those components coordinating and operating cores **2202A-N**. The system agent unit **2210** may include for example a power control unit (PCU) and a display unit. The PCU may be or include logic and components needed for regulating the power state of the cores **2202A-N** and the integrated graphics logic **2208**. The display unit is for driving one or more externally connected displays.

(249) The cores **2202A-N** may be homogenous or heterogeneous in terms of architecture instruction set; that is, two or more of the cores **2202A-N** may be capable of execution the same instruction set, while others may be capable of executing only a subset of that instruction set or a different instruction set.

(250) Exemplary Computer Architectures

(251) FIGS. **23-24** are block diagrams of exemplary computer architectures. Other system designs and configurations known in the arts for laptops, desktops, handheld PCs, personal digital assistants, engineering workstations, servers, network devices, network hubs, switches, embedded processors, digital signal processors (DSPs), graphics devices, video game devices, set-top boxes, micro controllers, cell phones, portable media players, hand held devices, and various other electronic devices, are also suitable. In general, a huge variety of systems or electronic devices capable of incorporating a processor and/or other execution logic as disclosed herein are generally suitable.

(252) Referring now to FIG. **23**, shown is a block diagram of a system **2300** in accordance with one embodiment of the present invention. The system **2300** may include one or more processors **2310**, **2315**, which are coupled to a controller hub **2320**. In one embodiment the controller hub **2320** includes a graphics memory controller hub (GMCH) **2390** and an Input/Output Hub (IOH) **2350** (which may be on separate chips); the GMCH **2390** includes memory and graphics controllers to which are coupled memory **2340** and a coprocessor **2345**; the IOH **2350** is couples input/output (I/O) devices **2360** to the GMCH **2390**. Alternatively, one or both of the memory and graphics controllers are integrated within the processor (as described herein), the memory **2340** and the coprocessor **2345** are coupled directly to the processor **2310**, and the controller hub **2320** in a single chip with the IOH **2350**.

(253) The optional nature of additional processors **2315** is denoted in FIG. **23** with broken lines. Each processor **2310**, **2315** may include one or more of the processing cores described herein and may be some version of the processor **2200**.

(254) The memory **2340** may be, for example, dynamic random access memory (DRAM), phase change memory (PCM), or a combination of the two. For at least one embodiment, the controller hub **2320** communicates with the processor(s) **2310**, **2315** via a multi-drop bus, such as a frontside bus (FSB), point-to-point interface such as QuickPath Interconnect (QPI), or similar connection **2395**.

(255) In one embodiment, the coprocessor **2345** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like. In one embodiment, controller hub **2320** may include an integrated graphics accelerator.

(256) There can be a variety of differences between the physical resources **2310**, **2315** in terms of a spectrum of metrics of merit including architectural, microarchitectural, thermal, power consumption characteristics, and the like.

(257) In one embodiment, the processor **2310** executes instructions that control data processing operations of a general type. Embedded within the instructions may be coprocessor instructions. The processor **2310** recognizes these coprocessor instructions as being of a type that should be executed by the attached coprocessor **2345**. Accordingly, the processor **2310** issues these coprocessor instructions (or control signals representing coprocessor instructions) on a coprocessor bus or other interconnect, to coprocessor **2345**. Coprocessor(s) **2345** accept and execute the received coprocessor instructions.

(258) Referring now to FIG. **24**, shown is a block diagram of a SoC **2400** in accordance with an embodiment of the present invention. Similar elements in FIG. **22** bear like reference numerals. Also, dashed lined boxes are optional features on more advanced SoCs. In FIG. **24**, an interconnect unit(s) **2402** is coupled to: an application processor **2410** which includes a set of one or more cores **202A-N** and shared cache unit(s) **2206**; a system agent unit **2210**; a bus controller unit(s) **2216**; an integrated memory controller unit(s) **2214**; a set of one or more coprocessors **2420** which may include integrated graphics logic, an image processor, an audio processor, and a video processor; a static random access memory (SRAM) unit **2430**; a direct memory access (DMA) unit **2432**; and a display unit **2440** for coupling to one or more external displays. In one embodiment, the coprocessor(s) **2420** include a special-purpose processor, such as, for example, a network or communication processor, compression engine, GPGPU, a high-throughput MIC processor, embedded processor, or the like.

(259) Embodiments of the mechanisms disclosed herein may be implemented in hardware, software, firmware, or a combination of such implementation approaches. Embodiments of the invention may be implemented as computer programs or program code executing on programmable systems comprising at least one processor, a storage system (including volatile and non-volatile memory and/or storage elements), at least one input device, and at least one output device.

(260) Program code may be applied to input instructions to perform the functions described herein and generate output information. The output information may be applied to one or more output devices, in known fashion. For purposes of this application, a processing system includes any system that has a processor, such as, for example; a digital signal processor (DSP), a microcontroller, an application specific integrated circuit (ASIC), or a microprocessor.

(261) The program code may be implemented in a high level procedural or object oriented programming language to communicate with a processing system. The program code may also be implemented in assembly or machine language, if desired. In fact, the mechanisms described herein are not limited in scope to any particular programming language. In any case, the language may be a compiled or interpreted language.

(262) One or more aspects of at least one embodiment may be implemented by representative instructions stored on a machine-readable medium which represents various logic within the processor, which when read by a machine causes the machine to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

(263) Such machine-readable storage media may include, without limitation, non-transitory, tangible arrangements of articles manufactured or formed by a machine or device, including storage media such as hard disks, any other type of disk including floppy disks, optical disks, compact disk read-only memories (CD-ROMs), compact disk rewritable's (CD-RWs), and magneto-optical disks, semiconductor devices such as read-only memories (ROMs), random access memories (RAMs) such as dynamic random access memories (DRAMs), static random access memories (SRAMs), erasable programmable read-only memories (EPROMs), flash memories, electrically erasable programmable read-only memories (EEPROMs), phase change memory (PCM), magnetic or optical cards, or any other type of media suitable for storing electronic instructions.

(264) Accordingly, embodiments of the invention also include non-transitory, tangible machine-readable media containing instructions or containing design data, such as Hardware Description Language (HDL), which defines structures, circuits, apparatuses, processors and/or system features described herein. Such embodiments may also be referred to as program products.

(265) Emulation (Including Binary Translation, Code Morphing, etc.)

(266) In some cases, an instruction converter may be used to convert an instruction from a source instruction set to a target instruction set. For example, the instruction converter may translate (e.g., using static binary translation, dynamic binary translation including dynamic compilation), morph, emulate, or otherwise convert an instruction to one or more other instructions to be processed by the core. The instruction converter may be implemented in software, hardware, firmware, or a combination thereof. The instruction converter may be on processor, off processor, or part on and part off processor.

(267) FIG. 25 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to embodiments of the invention. In the illustrated embodiment, the instruction converter is a software instruction converter, although alternatively the instruction converter may be implemented in software, firmware, hardware, or various combinations thereof. FIG. 25 shows a program in a high level language **2502** may be compiled using an x86 compiler **2504** to generate x86 binary code **2506** that may be natively executed by a processor with at least one x86 instruction set core **2516**. The processor with at least one x86 instruction set core **2516** represents any processor that can perform substantially the same functions as an Intel processor with at least one x86 instruction set core by compatibly executing or otherwise processing (1) a substantial portion of the instruction set of the Intel x86 instruction set core or (2) object code versions of applications or other software targeted to run on an Intel processor with at least one x86 instruction set core, in order to achieve substantially the same result as an Intel processor with at least one x86 instruction set core. The x86 compiler **2504** represents a compiler that is operable to generate x86 binary code **2506** (e.g., object code) that can, with or without additional linkage processing, be executed on the processor with at least one x86 instruction set core **2516**. Similarly, FIG. 25 shows the program in the high level language **2502** may be compiled using an alternative instruction set compiler **2508** to generate alternative instruction set binary code **2510** that may be natively executed by a processor without at least one x86 instruction set core **2514** (e.g., a processor with cores that execute the MIPS instruction set of MIPS Technologies of Sunnyvale, CA and/or that execute the ARM instruction set of ARM Holdings of Sunnyvale, CA). The instruction converter **2512** is used to convert the x86 binary code **2506** into code that may be natively executed by the processor without an x86 instruction set core **2514**. This converted code is not likely to be the same as the alternative instruction set binary code **2510** because an instruction converter capable of this is difficult to make; however, the converted code will accomplish the general operation and be made up of instructions from the alternative instruction set. Thus, the instruction converter **2512** represents software, firmware, hardware, or a combination thereof that, through emulation, simulation or any other process, allows a processor or other electronic device that does not have an x86 instruction set processor or core to execute the x86 binary code **2506**.

(268) Distribution of Frequency Budget in a Processor

(269) FIG. 26—Example System

(270) Referring now to FIG. 26, shown is a block diagram of a computing system **2600** in accordance with one or more embodiments. In some embodiments, the system **2600** may be all or a portion of an electronic device or component. For example, the system **2600** may be a cellular telephone, a computer, a server, a network device, a system on a chip (SoC), a controller, a wireless transceiver, a power supply unit, etc. Furthermore, in some embodiments, the system **2600** may be part of a grouping of related or interconnected devices, such as a datacenter, a computing cluster, etc.

(271) As shown in FIG. 26, the system 2600 may include a processor 2610 operatively coupled to system memory 2650. Further, although not shown in FIG. 26, the system 2600 may include other components. In one or more embodiments, the system memory 2650 can be implemented with any type(s) of computer memory (e.g., dynamic random-access memory (DRAM), static random-access memory (SRAM), non-volatile memory (NVM), a combination of DRAM and NVM, etc.).

(272) In one or more embodiments, the processor 2610 may be a hardware processing device (e.g., a central processing unit (CPU), a System on a Chip (SoC), and so forth). As shown, the processor 2610 can include any number of processing cores 2620A-2620N (also referred to generally as processing cores 2620, or alternatively as processing engines 2620), an internal network 2630, and a processor power management circuit (PMC) 2640. The processor PMC 2640 may include hardware (e.g., circuitry) and/or software to perform power management tasks for the processor 2610 and its components. The processing cores 2620 may include any type(s) of processing engine, including general-purpose processing engines, graphics processing engines, math co-processors, field programmable gate arrays (“FPGAs”), cryptographic processing engines, network processing engines, and so forth.

(273) As shown in FIG. 26, each processing core 2620 may include a core power management circuit (PMC) 2625. The core PMC 2625 may include hardware (e.g., circuitry) and/or software to perform power management tasks for its respective processing core 2620. In some embodiments, the core PMC 2625 may execute a prediction algorithm using core metrics as inputs. The prediction algorithm may generate a set of scalability hint values that indicate an amount of performance change that is associated with a frequency change to a given domain.

(274) In some embodiments, the processor PMC 2640 may determine total core gain and total uncore gain values based on the scalability hint values received from the processing cores 2620. The total core gain and total uncore gain values may be used to determine whether the processor is core bound, uncore bound, or memory bound. The processor PMC 2640 may then determine a frequency budget distribution based on total core gain and total uncore gain values. An example process for determining the frequency budget distribution is discussed below with reference to FIG. 28. In some embodiments, the frequency budget distribution generated for the processor 2610 may provide relatively efficient performance of heterogeneous workloads in various processing cores 2620. Further, the frequency budget distribution may be determined with a relatively short response time, and may therefore allow the processor to take advantage of dynamic workloads of short duration.

(275) FIG. 27—Example System

(276) Referring now to FIG. 27, shown is a system 2700 in accordance with some embodiments. As shown, the system 2700 may include any multiple engines 2740 interconnected by a mesh network 2720. In some embodiments, the engines 2740 may include general purpose processing engines (i.e., “cores”), graphics processing engines, math processing engines, network processing engines, cryptography processing engines, and so forth. The engines 2740 may correspond generally to the processing engine 2620 (shown in FIG. 26). Accordingly, the engines 2740 may represent all or a portion of a core domain of the processor 2610 (shown in FIG. 26), and may receive a first clock signal. Further, the mesh network 2720 may represent all or a portion of an uncore domain of the processor 2610, and may receive a second clock signal.

(277) In one or more embodiments, the mesh network 2720 may include multiple interconnected mesh stops 2730, where each mesh stop 2730 corresponds to an intersection or node of the mesh network 2720. Each mesh stop 2730 may include internal circuitry (not shown in FIG. 27) to provide various network-related functionality (e.g., routing circuits, cache memory, and so forth). In some embodiments, the mesh network 2720 may be implemented on a single die. Further, the mesh network 2720 may be connected to any number of other dies included in the system 2700 (not shown in FIG. 27). Note that, in some embodiments, the mesh network 2720 may be replaced with a Network on Chip (NoC), a fabric interconnect, and so forth.

(278) FIG. 28—Example Method

(279) Referring now to FIG. 28, shown is a flow diagram of a method **2800** for generating a frequency budget distribution, in accordance with one or more embodiments. In various embodiments, the method **2800** may be performed by processing logic (e.g., processor power management circuit (PMC) **2640** and/or core PMC **2625** shown in FIG. 26) that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, etc.), software (e.g., instructions run on a processing device), or a combination thereof. In firmware or software embodiments, the method **2800** may be implemented by computer executed instructions stored in a non-transitory machine-readable medium, such as an optical, semiconductor, or magnetic storage device. The machine-readable medium may store data, which if used by at least one machine, causes the at least one machine to fabricate at least one integrated circuit to perform a method.

(280) Block **2810** may include determining an average power use for a processor including multiple processing cores. Block **2815** may include determining power budgets for the processing cores. For example, referring to FIG. 26, telemetry circuitry may sample the voltage rails of the processor **2610** in a previous time cycle, and may generate metrics and/or register values indicating the energy consumed along with time stamps. The processor PMC **2640** may compute the power consumed during the previous cycle using these power metrics. Further, the processor PMC **2640** may apply a Running Average Power Limit (RAPL) algorithm to maintain a certain power limit averaged over a specific time window. The processor PMC **2640** may continuously monitor the power consumption, and may set limits on performance states of the processor **2610** to restrict the power consumption. Further, the processor PMC **2640** may perform power budgeting to the core, uncore, and/or memory domains of the processor **2610**. In some embodiments, the power budget may be computed using the formula $PB(n) = PB(n-1) + (1 - \exp(-DT/TW)) * (PL - DE/DT - PB(n-1))$. In this power budget formula, $PB(n)$ and $PB(n-1)$ are power budget of the current RAPL run and the last RAPL run, DT is delta time since the last RAPL run, DE is delta energy since the last RAPL run, DE/DT is average power consumed (e.g., determined at block **2810**) since the last RAPL run; TW is a programmed time window of the RAPL instance, and PL is a programmed power limit of the RAPL instance.

(281) Referring again to FIG. 28, block **2820** may include determining frequency limits for the processor. Block **2825** may include determining a core budget (CB) for the processor. Block **2830** may include determining an uncore budget (UB) for the processor. Block **2835** may include determining uncore and memory frequencies. The uncore frequency may be provided to an uncore domain **2880**, and the memory frequency may be provided to a memory domain **2890**. For example, referring to FIG. 26, a proportional integral derivative (PID) controller may be used to maintain a target budget by generating frequency limits for the core domain and the uncore domain. The power budget may be the input to the PID, and various parameters constants (e.g., K_p , K_i , and K_d parameters) may be tuned as needed. In some embodiments, there may be various power limits for the processor **2610**, and therefore multiple instances of power budget managers and PIDs may be run. Further, a minimum suggested frequency limit may be used as the core frequency ceiling limit. In the event that the power is constrained, the uncore frequency may be reduced along with the core frequency according to a linear function. In some examples, the uncore frequency ceiling limits may be derived from core limits using the formula $(\text{uncore frequency ceiling}) = (\text{minimum uncore frequency}) + ((\text{core frequency ceiling}) \times (\text{uncore frequency slope}))$. In some embodiments, the processor PMC **2640** may determine the uncore and memory frequencies based at least in part on workload (WL) hints, which may be generated at block **2864** (discussed below). Further, the processor PMC **2640** may determine the uncore and memory frequencies by executing an uncore frequency scaling (UFS) algorithm. The UFS algorithm may use multiple inputs including the WL hints, the maximum uncore bandwidth, a memory congestion indicator, an uncore congestion indicator, a memory bandwidth value, and so forth. An example process for using the WL hints to

adjust the uncore and memory frequencies is discussed below with reference to FIG. 33.

(282) Referring again to FIG. 28, block **2840** may include determining a per core budget. For example, referring to FIG. 26, the processor PMC **2640** may execute one or more algorithms to allocate the frequency budget to each processor core **2620**. For example, the processor PMC **2640** may execute a License Aware Running Average Power Limit (LA-RAPL) algorithm, including using different frequency limits for different licenses associated with instruction types (e.g., Streaming SIMD Extensions (SSE), Advanced Vector Extensions (AVX), and so forth). For example, different cores may choose a frequency ceiling limit based on the different AVX level instructions that are run on the core. In another example, the processor PMC **2640** may execute a Class of Service (COS) algorithm, including allocating a frequency ceiling limit to a group of cores having a particular priority. The group of cores may operate in a particular range of frequencies, with all cores with the same priority being allocated the same frequency ceiling limit.

(283) Referring again to FIG. 28, block **2845** may include determining a per core frequency (PCF). The per core frequency may be provided to the power management circuits (PMCs) **2855** included in the processing cores **2850**. For example, referring to FIG. 26, the processor PMC **2640** may execute a P-state selection algorithm, including using power budgets, per core scalability, and voltage and frequency levels to determine per core frequency limits. The per core frequency limits may be provided to the core PMC **2625** in each processing core **2620** included in the processor **2610**.

(284) Referring again to FIG. 28, in some embodiments, the PMCs **2855** may generate a set of scalability hint values indicating the performance change that may correspond to a frequency change in a given domain. For example, referring to FIG. 26, the core PMC **2625** may execute a prediction algorithm using core metrics as inputs, and may generate outputs including an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication. An example process for generating scalability hint values is discussed below with reference to FIG. 29.

(285) Referring again to FIG. 28, block **2860** may include determining core and uncore gain values for each processing core. For example, referring to FIG. 26, the processor PMC **2640** may calculate core and uncore gain values based on the scalability hint values received from the PMCs **2855**. The core gain may be a value indicating the expected gain in IPS performance for a specific step change of frequency in the core domain. Further, the uncore gain may indicate the expected gain in IPS performance for a frequency change in the uncore domain. Example formulas for calculating the core and uncore gain values are discussed below with reference to FIG. 30.

(286) Referring again to FIG. 28, block **2862** may include determining the total core gain and the total uncore gain for the processor. For example, referring to FIG. 26, the processor PMC **2640** may sum the core gain values for all processing cores **2620** to determine the total core gain for the processor (i.e., the expected gain in IPS performance across the processor for a frequency change in the core domain). Further, the processor PMC **2640** may sum all of the uncore gain values to determine the total uncore gain for the processor.

(287) Referring again to FIG. 28, block **2864** may include generating a set of workload (WL) hints. The WL hints may be provided as inputs for determining the uncore and memory frequencies at block **2835** (discussed above). For example, referring to FIG. 26, the processor PMC **2640** may generate WL hints that indicate the workload profile across all processing cores **2620**. The WL hints may indicate that a processing core **2620** is core bound, uncore bound, memory bound, or is at an ideal balance between domains. Further, the WL hints may indicate that the processing core **2620** is quiesced, or has balanced uncore and memory domains. In some embodiments, the WL hints may be determined based on inputs that may include the total core gain value and the total uncore gain value (determined at block **2862**), a cache miss rate, a memory bandwidth value, a memory congestion value, and an uncore bandwidth value. An example process for determining the WL hints is discussed below with reference to FIG. 33.

(288) Referring again to FIG. 28, block 2866 may include determining a distribution adjustment value based on the total core gain and total uncore gain (determined at block 2862). Block 2870 may include determining an updated budget distribution based on the distribution adjustment value. Block 2872 may include adjusting the core budget (CB) based on the updated budget distribution. Block 2874 may include adjusting the uncore budget (UB) based on the updated budget distribution. For example, referring to FIG. 26, the processor PMC 2640 may compute a ratio of the total core gain to the total uncore gain (also referred to herein as the “C/U gain ratio”) by dividing the total core gain by the total uncore gain. The processor PMC 2640 may use the C/U gain ratio to determine the amount of increase or decrease for the core and uncore frequency budget limits. For example, a C/U gain ratio greater than one may indicate that the workload is core bound, and therefore may indicate that more frequency budget should be distributed to the core domain. In another example, a C/U gain ratio less than one may indicate that the workload is uncore bound, and therefore may indicate that more frequency budget should be distributed to the uncore domain. The processor PMC 2640 may then modify the core and uncore frequencies based on the budget distribution. The method 2800 may then return to block 2840 to form a loop, which may be repeated to continuously adjust the frequency budget distribution of the processor. In this manner, the method 2800 may provide a budget distribution that may result in relatively efficient performance of heterogeneous workloads, and that may take advantage of dynamic workloads of relatively short duration.

(289) In some embodiments, the determination of the budget distribution may also be based on an uncore utilization value. For example, if the uncore utilization value indicates that the input/output traffic on a path from a mesh fabric to the memory, a decision to decrease the uncore frequency (e.g., based on a core bound indication) may be overridden to maintain the current frequency in the uncore domain. Further, in some embodiments, the amount of frequency change (i.e., increase or decrease) may be defined in terms of frequency bins. As used herein, the term “bins” refers to equal subdivisions of a frequency range. For example, the new core frequency limit may be calculated by the formula $\text{New_Core_Limit} = \text{Old_Core_Limit} \pm M \text{ bins}$. Further, the new uncore frequency limit may be calculated by the formula $\text{New_Uncore_Limit} = \text{Old_Uncore_Limit} \pm N \text{ bins}$. In this formula, the variables M and N may be integers that indicate the number of bins to be added or subtracted. The M and N variables may be determined using a bin ratio that translates changes in core frequency to changes in uncore frequency according to equivalent power consumption levels. The bin ratio may vary according to different frequency levels. For example, at a first frequency level, an increase of one bin (e.g., 100 MHz) in the core frequency may represent an equivalent power level to an increase of two bins (e.g., 200 Mhz) in the uncore frequency. In another example, at a second frequency level, an increase of one bin (e.g., 100 MHz) in the core frequency may represent an equivalent power level to an increase of three bins (e.g., 300 Mhz) in the uncore frequency. In some embodiments, the bin ratios associated with different frequencies may be obtained from stored data structures (e.g., stored tables), and may be based on observed measurements and/or user configuration settings. An example process for determining the core and uncore frequency limits is described below with reference to FIGS. 31-32.

(290) Referring again to FIG. 28, in some embodiments, determining a per core budget (at block 2840) may be based at least in part on the core scalability (CS) values generated by the PMCs 2855 (discussed above). For example, as discussed above with reference to block 2840, the processor PMC 2640 may use the core budget (CB) determined at block 2825 (e.g., using a per core frequency distribution algorithm or P-state stack) to increase or decrease the core frequency in line with the workload phase behavior. Further, the processor PMC 2640 (e.g., executing a p-stack) may use the core scalability (CS) values to determine how to distribute the core budget to the various processing cores. In this manner, the processor PMC 2640 may receive feedback on the core scalability via the CS values for each processing core.

(291) FIG. 29—Example Operation

(292) Referring now to FIG. 29, shown is an example operation 2900 for generating core scalability hints, in accordance with one or more embodiments. As shown in FIG. 29, a core power management circuit (PMC) 2625 may execute a prediction algorithm 2910 to generate a set of core scalability hints 2920. The prediction algorithm 2910 may use various core metrics as inputs. The core metrics may indicate characteristics of the workload in a respective core in a recent period of time. For example, the core PMC 2625 may execute a prediction algorithm that uses various core telemetry counters associated with the workload of the respective processing core 2620.

(293) In some embodiments, the core telemetry counters used by the prediction algorithm 2910 may include the instructions per cycle (IPC), the number of inactive cycles, counts of queue occupancy for multiple thresholds, the number of prefetches in a cycle and whether the prefetch was effective, the number of cache hits, the number of cache misses, stalls for caches at levels 1, 2, and 3 (i.e., L1/L2/L3 stalls), the number of cycles that the core was active (e.g., in a C0 state), the number of cycles that the core was inactive, the average frequency for the core during the workload, and any other core metric.

(294) In some embodiments, the core scalability hints 2920 may include values that indicate an amount of performance change that is associated with a frequency change to a given domain. For example, the core scalability hints 2920 may include an instructions per second (IPS) value 2922, a core scalability (CS) value 2924, an uncore scalability (US) value 2926, and a quiesced core indication 2928. The IPS value 2922 may indicate the current performance level of the processing core 2620 (shown in FIG. 26), and may be calculated by multiplying an instructions per cycle (IPC) metric times the current clock frequency of the processing core 2620.

(295) In some embodiments, the core scalability value 2924 may indicate the expected change in performance for a given change of frequency in the core domain, and may be calculated as the ratio between the IPS value 2922 and the frequency change in the core domain. The uncore scalability value 2926 may indicate the expected change in performance for a given change of frequency in the uncore domain, and may be calculated as the ratio between the IPS value 2922 and the frequency change in the uncore domain. The quiesced core indication 2928 may indicate that the respective processing core 2620 is inactive. In some examples, the core scalability hints 2920 may vary according to factors such as the number of memory dependency periods (e.g., periods where the processing core 2620 waits for data to return from the uncore, and does not execute any instructions), the number of pause loops, and so forth.

(296) FIG. 30—Example Formulas

(297) Referring now to FIG. 30, shown are example formulas used in accordance with one or more embodiments. For example, the first formula 3000 may be used for calculating the scalability values (discussed above with reference to FIG. 29). In the first formula 3000, the term “Block” may represent either “core” or “uncore” for each calculation. In another example, the second formula 3010 may be used for calculating the core gain values (discussed above with reference to block 2860 of FIG. 28). The core gain may be a value indicating the expected gain in IPS performance for a given change of frequency in the core domain, and may be calculated as the product of the core scalability value, times the IPS value, times the bin ratio. In yet another example, the third formula 3020 may be used for calculating the uncore gain values (discussed above with reference to block 2860 of FIG. 28). The uncore gain may be a value indicating the expected gain in IPS performance for a given change of frequency in the uncore domain, and may be calculated as the product of the uncore scalability value, times the IPS value, times the bin ratio.

(298) FIG. 31—Example Table

(299) Referring now to FIG. 31, shown is an example core bound table 3100, in accordance with one or more embodiments. The core bound table 3100 may illustrate the frequency adjustments made under a core bound scenario. For example, the core bound table 3100 may be implemented in logic or stored data, and may be used to determine the distribution adjustment value based on the ratio of the total core gain to the total uncore gain (“C/U gain ratio”), as discussed above with

reference to block **2866** (shown in FIG. **28**).

(300) As shown in FIG. **31**, the core bound table **3100** may include three entries, corresponding to low, medium, and high threshold levels of the C/U gain ratio. For example, the low level may correspond to a C/U gain ratio between 20% and 40%, the medium level may correspond to a C/U gain ratio between 40% and 70%, and the high level may correspond to a C/U gain ratio above 70%. Further, each entry may specify a number M of bins to increase in the core frequency, and a number N of bins to decrease in the uncore frequency.

(301) FIG. **32**—Example Table

(302) Referring now to FIG. **32**, shown is an example uncore bound table **3200**, in accordance with one or more embodiments. The uncore bound table **3200** may illustrate the frequency adjustments made under an uncore bound scenario. The core bound table **3200** may be implemented in logic or stored data. As shown, the uncore bound table **3200** may include three entries, corresponding to low, medium, and high threshold levels of the C/U gain ratio. Further, each entry may specify a number M of bins to decrease in the core frequency, and a number N of bins to increase in the uncore frequency.

(303) FIG. **33**—Example Table

(304) Referring now to FIG. **33**, shown is an example hint table **3300**, in accordance with one or more embodiments. The hint table **3300** may illustrate the workload (WL) hints that may be generated at block **2864** based on the total core gain and total uncore gain values. Further, the hint table **3300** may illustrate the changes to core, uncore, and memory frequencies (e.g., performed in block **2835** of FIG. **28**) that may be associated with the various hints. For example, referring to FIG. **26**, the processor PMC **2640** may generate hint **H0** in response to a determination that the total core gain is larger than the total uncore gain, indicating that the workload is core bound. Accordingly, the hint **H0** may be associated with an increase to the core frequency, and a reduction to the uncore frequency and the memory frequency.

(305) The hint **H1** may be generated in response to a determination that the total core gain is equal to the total uncore gain, indicating an ideal balance between these domain frequencies in the processing core. Accordingly, the hint **H1** may be associated with no changes to the core frequency, the uncore frequency, and the memory frequency.

(306) The hint **H2** may be generated in response to a determination that the workload of the processing core is uncore bound (e.g., the total uncore gain is larger than the total core gain, and that the workload is not memory bound). Accordingly, the hint **H2** may be associated with a reduction to the core frequency, an increase to the uncore frequency, and a reduction to the memory frequency.

(307) The hint **H3** may be generated in response to a determination that the workload of the processing core is memory bound (e.g., the total uncore gain is larger than the total core gain, and one of the following is true: there is a high cache miss rate, there is high memory bandwidth, or there is high memory congestion). Accordingly, the hint **H3** may be associated with a reduction to the core frequency, a reduction to the uncore frequency, and an increase to the memory frequency.

(308) The hint **H4** may be generated in response to a determination that all (or more than a threshold number) of the processing cores are idle or clock-gated (e.g., as indicated by the quiesced core value). Accordingly, the hint **H4** may be associated with a reduction to the core frequency, a reduction to the uncore frequency, and a reduction to the memory frequency.

(309) The hint **H5** may be generated in response to a determination that the workload of the processing core is equally distributed between uncore bound and memory bound. Accordingly, the hint **H3** may be associated with no change to the core frequency, an increase to the uncore frequency, and an increase to the memory frequency.

(310) FIG. **34**—Example Method

(311) Referring now to FIG. **34**, shown is a flow diagram of a method **3400** for determining a frequency budget, in accordance with one or more embodiments. In various embodiments, the

method **3400** may be performed by processing logic (e.g., processor power management circuit (PMC) **2640** and/or core PMC **2625** shown in FIG. **26**) that may include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, etc.), software (e.g., instructions run on a processing device), or a combination thereof. In firmware or software embodiments, the method **3400** may be implemented by computer executed instructions stored in a non-transitory machine-readable medium, such as an optical, semiconductor, or magnetic storage device. The machine-readable medium may store data, which if used by at least one machine, causes the at least one machine to fabricate at least one integrated circuit to perform a method.

(312) Block **3410** may include determining, for each processing core of a processor, a plurality of metrics indicating a current workload of the processing core, wherein the processor comprises a core domain and an uncore domain, and wherein the core domain comprises a plurality of processing cores. Block **3420** may include executing by a power management circuit of the processor, using the plurality of metrics as inputs, a prediction algorithm to generate a plurality of scalability hint values. For example, referring to FIGS. **26-29**, each processing core **2620** may include a core power management circuit (PMC) **2625**. Each core PMC **2625** may execute a prediction algorithm **2910** to generate a set of core scalability hints **2920** for the corresponding processing core **2620**. In some embodiments, the prediction algorithm **2910** may use various core telemetry counters associated with the workload of the corresponding processing core **2620**.

(313) Block **3430** may include determining, by the power management circuit, a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values. For example, referring to FIG. **26**, the processor PMC **2640** may calculate core and uncore gain values based on the scalability hint values received from the core PMCs **2625**. The processor PMC **2640** may sum the core gain values for all processing cores **2620** to determine the total core gain for the processor. Further, the processor PMC **2640** may sum all of the uncore gain values to determine the total uncore gain for the processor.

(314) Block **3440** may include distributing, by the power management circuit, a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. After block **3440**, the method **3400** may be completed. For example, referring to FIG. **26**, the processor PMC **2640** may compute a ratio of the total core gain to the total uncore gain (“C/U gain ratio”) by dividing the total core gain by the total uncore gain. The processor PMC **2640** may use the C/U gain ratio to determine the amount of increase or decrease for the core and uncore frequency budget limits. The processor PMC **2640** may then modify the core and uncore frequencies based on the budget distribution. In some embodiments, the method **3400** may be repeated periodically to continuously adjust the frequency budget distribution of the processor. In this manner, the method **3400** may provide a budget distribution that may result in relatively efficient performance of heterogeneous workloads, and that may take advantage of dynamic workloads of relatively short duration.

(315) FIG. **35**—Example Storage Medium

(316) Referring now to FIG. **35**, shown is a storage medium **3500** storing executable instructions **3510**. In some embodiments, the storage medium **3500** may be a non-transitory machine-readable medium, such as an optical medium, a semiconductor, a magnetic storage device, and so forth. The executable instructions **3510** may be executable by a processing device. Further, the executable instructions **3510** may be used by at least one machine to fabricate at least one integrated circuit to perform one or more of the methods and/or operations shown in FIGS. **28-34**.

(317) The following clauses and/or examples pertain to further embodiments. In Example 1, a processor for power management may include a core domain comprising a plurality of processing cores; an uncore domain comprising an internal network; and a processor power management circuit. The processor power management circuit may be to: receive a plurality of scalability hint values from the plurality of processing cores; determine a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distribute a frequency budget

between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. In Example 2, the subject matter of Example 1 may optionally include that the processing core includes a core power management circuit to: determine a plurality of metrics that indicate a current workload of the processing core; and execute, using the plurality of metrics as inputs, a prediction algorithm to generate the plurality of scalability hint values. In Example 3, the subject matter of Examples 1-2 may optionally include that the plurality of metrics comprises instructions per cycle, count of inactive cycles, counts of queue occupancy for multiple thresholds, count of prefetches in a cycle, count of cache hits, and count of cache misses. In Example 4, the subject matter of Examples 1-3 may optionally include that the internal network comprises a mesh network including a plurality of mesh stops. In Example 5, the subject matter of Examples 1-4 may optionally include that the plurality of scalability hint values comprises: an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication. In Example 6, the subject matter of Examples 1-5 may optionally include that processor power management circuit is to: for each of the plurality of processing cores, calculate a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculate an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio. In Example 7, the subject matter of Examples 1-6 may optionally include that the processor power management circuit is to: calculate the total core gain as a sum of the core gain values for the plurality of processing cores; and calculate the total uncore gain as a sum of the uncore gain values for the plurality of processing cores. In Example 8, the subject matter of Examples 1-7 may optionally include that the processor power management circuit is to: generate, based at least on the total core gain and the total uncore gain, a workload hint to indicate a workload profile of the processing core; and determine an uncore frequency based at least on the generated workload hint. In Example 9, the subject matter of Examples 1-8 may optionally include that the processor power management circuit is to determine a per core frequency budget based at least in part on the core scalability values for each of the plurality of processing cores. In Example 10, a method for power management may include: determining, for each processing core of a processor, a plurality of metrics indicating a current workload of the processing core, where the processor comprises a core domain and an uncore domain, and wherein the core domain comprises a plurality of processing cores; executing by a power management circuit of the processor, using the plurality of metrics as inputs, a prediction algorithm to generate a plurality of scalability hint values; determining, by the power management circuit, a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distributing, by the power management circuit, a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. In Example 11, the subject matter of Example 10 may optionally include that the plurality of scalability hint values comprises: an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication. In Example 12, the subject matter of Examples 10-11 may optionally include: for each of the plurality of processing cores, calculating a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculating an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio. In Example 13, the subject matter of Examples 10-12 may optionally include: calculating the total core gain as a sum of the core gain values for the plurality of processing cores; and calculating the total uncore gain as a sum of the uncore gain values for the plurality of processing cores. In Example 14, the subject matter of Examples 10-13 may optionally include determining a per core frequency budget based at least in part on the core scalability values for each of the plurality of processing cores. In Example 15, the subject matter of Examples 10-14 may optionally include: generating, based at least on the total core gain and the total uncore gain, a workload hint to indicate a workload profile of the processing core; and determining an uncore frequency based at least on the generated workload hint. In

Example 16, the subject matter of Examples 10-15 may optionally include: computing a ratio of the total core gain to the total uncore gain; determining a budget distribution based on the ratio of the total core gain to the total uncore gain; and distributing a frequency budget based on the determined budget distribution. In Example 17, a computing device may include one or more processors, and a memory having stored therein a plurality of instructions that when executed by the one or more processors, cause the computing device to perform the method of any of Examples 10 to 16. In Example 18, a machine readable medium may have stored thereon data, which if used by at least one machine, causes the at least one machine to fabricate at least one integrated circuit to perform a method according to any one of Examples 10 to 16. In Example 19, an electronic device may include means for performing the method of any of Examples 10 to 16. In Example 20, a system for power management may include a processor, and a memory coupled to the processor. The processor may include a core domain comprising a plurality of processing cores, an uncore domain comprising an internal network, and a processor power management circuit. The processor power management circuit may be to: receive a plurality of scalability hint values from the plurality of processing cores; determine a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distribute a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. In Example 21, the subject matter of Example 20 may optionally include that the plurality of scalability hint values comprises an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication. In Example 22, the subject matter of Examples 20-21 may optionally include that the processor power management circuit is to: for each of the plurality of processing cores, calculate a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculate an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio. In Example 23, the subject matter of Examples 20-22 may optionally include that the processor power management circuit is to: calculate the total core gain as a sum of the core gain values for the plurality of processing cores; and calculate the total uncore gain as a sum of the uncore gain values for the plurality of processing cores. In Example 24, an apparatus for power management may include: means for determining, for each processing core of a processor, a plurality of metrics indicating a current workload of the processing core, where the processor comprises a core domain and an uncore domain, and wherein the core domain comprises a plurality of processing cores; means for executing, using the plurality of metrics as inputs, a prediction algorithm to generate a plurality of scalability hint values; means for determining a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and means for distributing a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain. In Example 25, the subject matter of Example 24 may optionally include that the plurality of scalability hint values comprises: an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication. In Example 26, the subject matter of Examples 24-25 may optionally include means for: for each of the plurality of processing cores, calculating a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculating an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio. In Example 27, the subject matter of Examples 24-26 may optionally include: means for calculating the total core gain as a sum of the core gain values for the plurality of processing cores; and means for calculating the total uncore gain as a sum of the uncore gain values for the plurality of processing cores. In Example 28, the subject matter of Examples 24-27 may optionally include means for determining a per core frequency budget based at least in part on the core scalability values for each of the plurality of processing cores. In Example 29, the subject matter of Examples 24-28 may optionally include: means for generating, based at least on the total core gain and the total uncore gain, a workload hint to indicate a workload profile of the processing

core; and means for determining an uncore frequency based at least on the generated workload hint. In Example 30, the subject matter of Examples 24-29 may optionally include: means for computing a ratio of the total core gain to the total uncore gain; means for determining a budget distribution based on the ratio of the total core gain to the total uncore gain; and means for distributing a frequency budget based on the determined budget distribution.

(318) In accordance with one or more embodiments, a processor may include circuitry to determine a frequency budget distribution based on current workloads in the processor. The circuitry may receive scalability hint values from each processing core in the processor. The scalability hint values may be generated by a prediction algorithm based on multiple metrics that indicate the current workload of the respective core. The circuitry may determine total core gain and total uncore gain values based at least in part on the scalability hint values received from the cores. The total core gain and total uncore gain values may be used to determine whether the processor is core bound, uncore bound, or memory bound. The circuitry may determine the frequency budget distribution based at least in part on total core gain and total uncore gain values. In this manner, some embodiments may provide a frequency budget distribution that improves the performance of heterogeneous workloads of the cores. Further, some embodiments may allow the frequency budget distribution to be determined with a relatively short response time, and may therefore allow the processor to take advantage of dynamic workloads of short duration.

(319) Note that, while FIGS. 26-35 illustrate various example implementations, other variations are possible. For example, it is contemplated that one or more embodiments may be implemented in the example devices and systems described with reference to FIGS. 1-25. Note that the examples shown in FIGS. 1-35 are provided for the sake of illustration, and are not intended to limit any embodiments. Specifically, while embodiments may be shown in simplified form for the sake of clarity, embodiments may include any number and/or arrangement of components. For example, it is contemplated that some embodiments may include any number of components in addition to those shown, and that different arrangement of the components shown may occur in certain implementations. Furthermore, it is contemplated that specifics in the examples shown in FIGS. 1-35 may be used anywhere in one or more embodiments.

(320) Understand that various combinations of the above examples are possible. Embodiments may be used in many different types of systems. For example, in one embodiment a communication device can be arranged to perform the various methods and techniques described herein. Of course, the scope of the present invention is not limited to a communication device, and instead other embodiments can be directed to other types of apparatus for processing instructions, or one or more machine readable media including instructions that in response to being executed on a computing device, cause the device to carry out one or more of the methods and techniques described herein.

(321) References throughout this specification to “one embodiment” or “an embodiment” mean that a particular feature, structure, or characteristic described in connection with the embodiment is included in at least one implementation encompassed within the present invention. Thus, appearances of the phrase “one embodiment” or “in an embodiment” are not necessarily referring to the same embodiment. Furthermore, the particular features, structures, or characteristics may be instituted in other suitable forms other than the particular embodiment illustrated and all such forms may be encompassed within the claims of the present application.

(322) While the present invention has been described with respect to a limited number of embodiments, those skilled in the art will appreciate numerous modifications and variations therefrom. It is intended that the appended claims cover all such modifications and variations as fall within the true spirit and scope of this present invention.

Claims

1. A processor comprising: a core domain comprising a plurality of processing cores; an uncore domain comprising an internal network; and a processor power management circuit to: receive a plurality of scalability hint values from the plurality of processing cores; determine a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distribute a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain.
2. The processor of claim 1, wherein each processing core comprises: a core power management circuit to: determine a plurality of metrics that indicate a current workload of the processing core; and execute, using the plurality of metrics as inputs, a prediction algorithm to generate the plurality of scalability hint values.
3. The processor of claim 2, wherein the plurality of metrics comprises: instructions per cycle, count of inactive cycles, counts of queue occupancy for multiple thresholds, count of prefetches in a cycle, count of cache hits, and count of cache misses.
4. The processor of claim 1, wherein the internal network comprises a mesh network including a plurality of mesh stops.
5. The processor of claim 1, wherein the plurality of scalability hint values comprises: an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication.
6. The processor of claim 5, wherein the processor power management circuit is to: for each of the plurality of processing cores, calculate a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculate an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio.
7. The processor of claim 6, wherein the processor power management circuit is to: calculate the total core gain as a sum of the core gain values for the plurality of processing cores; and calculate the total uncore gain as a sum of the uncore gain values for the plurality of processing cores.
8. The processor of claim 6, wherein the processor power management circuit is to: generate, based at least on the total core gain and the total uncore gain, a workload hint to indicate a workload profile of the processing core; and determine an uncore frequency based at least on the generated workload hint.
9. The processor of claim 5, wherein the processor power management circuit is to: determine a per core frequency budget based at least in part on the core scalability values for each of the plurality of processing cores.
10. A method comprising: determining, for each processing core of a processor, a plurality of metrics indicating a current workload of the processing core, wherein the processor comprises a core domain and an uncore domain, and wherein the core domain comprises a plurality of processing cores; executing by a power management circuit of the processor, using the plurality of metrics as inputs, a prediction algorithm to generate a plurality of scalability hint values; determining, by the power management circuit, a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distributing, by the power management circuit, a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain.
11. The method of claim 10, wherein the plurality of scalability hint values comprises: an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication.
12. The method of claim 11, further comprising: for each of the plurality of processing cores, calculating a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculating an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio.
13. The method of claim 12, further comprising: calculating the total core gain as a sum of the core

gain values for the plurality of processing cores; and calculating the total uncore gain as a sum of the uncore gain values for the plurality of processing cores.

14. The method of claim 11, further comprising: determining a per core frequency budget based at least in part on the core scalability values for each of the plurality of processing cores.

15. The method of claim 10, further comprising: generating, based at least on the total core gain and the total uncore gain, a workload hint to indicate a workload profile of the processing core; and determining an uncore frequency based at least on the generated workload hint.

16. The method of claim 10, further comprising: computing a ratio of the total core gain to the total uncore gain; determining a budget distribution based on the ratio of the total core gain to the total uncore gain; and distributing a frequency budget based on the determined budget distribution.

17. A system comprising: a processor comprising: a core domain comprising a plurality of processing cores; an uncore domain comprising an internal network; and a processor power management circuit to: receive a plurality of scalability hint values from the plurality of processing cores; determine a total core gain and a total uncore gain based at least in part on the plurality of scalability hint values; and distribute a frequency budget between the core domain and the uncore domain based at least in part on the total core gain and the total uncore gain; and a memory coupled to the processor.

18. The system of claim 17, wherein the plurality of scalability hint values comprises an instructions per second (IPS) value, a core scalability (CS) value, an uncore scalability (US) value, and a quiesced core (QC) indication.

19. The system of claim 18, wherein the processor power management circuit is to: for each of the plurality of processing cores, calculate a core gain value by multiplying the core scalability value, the IPS value, and a bin ratio; and for each of the plurality of processing cores, calculate an uncore gain value by multiplying the uncore scalability value, the IPS value, and the bin ratio.

20. The system of claim 19, wherein the processor power management circuit is to: calculate the total core gain as a sum of the core gain values for the plurality of processing cores; and calculate the total uncore gain as a sum of the uncore gain values for the plurality of processing cores.
