

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417038
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Basit; Abdul et al.

Block allocation for persistent memory during aggregate transition

Abstract

Techniques are provided for block allocation for persistent memory during aggregate transition. In a high availability pair including first and second nodes, the first node makes a determination that control of a first aggregate is to transition from the first node to the second node. A portion of available free storage space is allocated from a first persistent memory of the first node as allocated pages within the first persistent memory. Metadata information for the allocated pages is updated with an identifier of the first aggregate to create updated metadata information reserving the allocated pages for the first aggregate. The updated metadata information is mirrored to the second node, so that the second node also reserves those pages. Control of the first aggregate is transitioned to the second node. As a result, the nodes do not attempt allocating the same free pages to different aggregates during a transition.

Inventors: Basit; Abdul (Morrisville, NC), Subramanian; Ananthan (San Ramon, CA), Kesavan; Ram (Los Altos, CA), Curtis-Maury; Matthew Fontaine (Apex, NC)

Applicant: NetApp Inc. (San Jose, CA)

Family ID: 1000008818605

Assignee: NetApp, Inc. (San Jose, CA)

Appl. No.: 18/528556

Filed: December 04, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240103744 A1	Mar. 28, 2024

Related U.S. Application Data

continuation parent-doc US 17750793 20220523 US 11836363 child-doc US 18528556

Publication Classification

Int. Cl.: **G06F3/06** (20060101)

U.S. Cl.:

CPC **G06F3/0631** (20130101); **G06F3/0604** (20130101); **G06F3/064** (20130101); **G06F3/065** (20130101); **G06F3/0679** (20130101);

Field of Classification Search

CPC: G06F (3/0631); G06F (3/0604); G06F (3/064); G06F (3/065); G06F (3/0679)

USPC: 711/162; 711/165

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10394485	12/2018	Chen et al.	N/A	N/A
10459806	12/2018	He et al.	N/A	N/A
11036677	12/2020	Grunwald et al.	N/A	N/A
11340804	12/2021	Basit et al.	N/A	N/A
11836363	12/2022	Basit et al.	N/A	N/A
2011/0314219	12/2010	Ulrich et al.	N/A	N/A
2012/0303685	12/2011	Ulrich et al.	N/A	N/A
2017/0132091	12/2016	Leshinsky	N/A	G06F 3/067
2017/0177448	12/2016	Muth	N/A	G06F 3/064
2018/0046552	12/2017	Madhavarapu	N/A	G06F 3/065
2018/0067662	12/2017	Leshinsky	N/A	G06F 16/219
2018/0217756	12/2017	Liu et al.	N/A	N/A
2020/0409583	12/2019	Kusters et al.	N/A	N/A
2022/0283718	12/2021	Basit et al.	N/A	N/A

OTHER PUBLICATIONS

Non final Office Action mailed Mar. 2, 2023 for U.S. Appl. No. 17/750,793, filed May 23, 2022, 9 pages. cited by applicant
 Notice of Allowance mailed on Jul. 28, 2023 for U.S. Appl. No. 17/750,793, filed May 23, 2022, 7 pages. cited by applicant

Primary Examiner: Nguyen; Than

Attorney, Agent or Firm: Cooper Legal Group, LLC

Background/Summary

RELATED APPLICATION (1) This application claims priority to and is a continuation of U.S. patent application Ser. No. 17/750,793, titled BLOCK ALLOCATION FOR PERSISTENT MEMORY DURING AGGREGATE TRANSITION” and filed on May 23, 2022, which priority to and is a continuation of U.S. Pat. No. 11,340,804, titled BLOCK ALLOCATION FOR PERSISTENT MEMORY DURING AGGREGATE TRANSITION” and filed on Jun. 25, 2020, which are incorporated herein by reference.

BACKGROUND

(1) A computing environment may host one or more nodes, such as servers, virtual machines, computing devices, etc., for storing data on behalf of clients. The nodes may be deployed in a manner that provides high availability, data redundancy, and/or other storage features. For example, a first node may host one or more aggregates within which data is stored, such as a first aggregate used to store data on behalf of a first client. The first node may store at least some data of the first aggregate within a first persistent memory of the first node. A second node may also host one or more aggregates within which data is stored, such as a second aggregate used to store data on behalf of a second client. The second node may store at least some data of the second aggregate within a second persistent memory of the second node.

(2) The first node and the second node may be configured as a node pairing (e.g., a high availability node pairing) configured to provide high availability and/or failover functionality. For example, if one of the nodes fails, then the surviving node can provide clients with failover access to their data in place of the failed node. This can be accomplished by mirroring data between the persistent memories of the nodes so that the surviving node has an up-to-date copy of data from a persistent memory of the failed node. In this way, the surviving node can provide clients with up-to-date data that was previously accessible to the clients through the failed node.

Description

DESCRIPTION OF THE DRAWINGS

(1) FIG. 1 is a block diagram illustrating an example computing environment in which an embodiment may be implemented.

(2) FIG. 2 is a block diagram illustrating a network environment with exemplary node computing devices.

(3) FIG. 3 is a block diagram illustrating an exemplary node computing device.

(4) FIG. 4 is a flow chart illustrating an example method for block allocation for persistent memory during aggregate transition.

(5) FIG. 5A is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (A) and a node (B) are operating under normal operation.

(6) FIG. 5B is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (B) fails.

(7) FIG. 5C is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (A) takes over an aggregate (B1) and an aggregate (B2).

(8) FIG. 5D is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (B) recovers from a failure and a node (A) updates metadata information for an aggregate (B1).

(9) FIG. 5E is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (A) gives back an aggregate (B1) to a node (B).

(10) FIG. 5F is a block diagram illustrating an example system for block allocation for persistent

memory during aggregate transition, where a node (A) updates metadata information for an aggregate (B2).

(11) FIG. 5G is a block diagram illustrating an example system for block allocation for persistent memory during aggregate transition, where a node (A) gives back an aggregate (B2) to a node (B).

(12) FIG. 6 is an example of a computer readable medium in which an embodiment may be implemented.

DETAILED DESCRIPTION

(13) Some examples of the claimed subject matter are now described with reference to the drawings, where like reference numerals are generally used to refer to like elements throughout. In the following description, for purposes of explanation, numerous specific details are set forth in order to provide an understanding of the claimed subject matter. It may be evident, however, that the claimed subject matter may be practiced without these specific details. Nothing in this detailed description is admitted as prior art.

(14) A first node may be implemented as a computing device, a server, an on-premise device, a virtual machine, hardware, software, or combination thereof. The first node may store client data within one or more aggregates, such as a first aggregate and a second aggregate. In order to provide high availability and redundancy, the first node may be paired with a second node. The second node may also store client data within one or more aggregates, such as a third aggregate and a fourth aggregate. Data is mirrored between the first node and the second node, such as between a first persistent memory of the first node and a second persistent memory of the second node, so that both nodes have up-to-date data of the other node.

(15) For example, the first persistent memory may be partitioned into a first local partition into which data of the first aggregate and the second aggregate may be stored. The first persistent memory may be partitioned into a first remote partition into which data of the third aggregate and the fourth aggregate is mirrored from the second persistent memory of the second node (e.g., from a second local partition of the second persistent memory). Similarly, the second persistent memory may be partitioned into the second local partition into which data of the third aggregate and the fourth aggregate may be stored. The second persistent memory may be partitioned into a second remote partition into which data of the first aggregate and second aggregate is mirrored from the first local partition of the first persistent memory of the first node.

(16) Because data is mirrored between the persistent memories of the nodes, the nodes may be capable of providing takeover and giveback functionality. For example, if one of the nodes fails, then the surviving node can take over aggregates of the failed node using up-to-date data within the remote partition of the persistent memory of the surviving node. Once the failed node has recovered, the surviving node can give back control of the taken over aggregates to the recovered node.

(17) For example, the first node may fail, and the second node may take over the first and second aggregates from the first node. The second node may scan the second remote partition of the second persistent memory, corresponding to the mirrored data of the first and second aggregates, to build a free pages list of free pages within the second remote partition that can be subsequently allocated to store data directed to the first and second aggregates. In this way, the second node will service I/O directed to the first and second aggregates using the second remote partition. Once the first node recovers from the failure, the second node may perform a giveback procedure to return the first and second aggregates to the first node one at a time. Also, once the first node has recovered, a resynchronization procedure is performed to copy data of the second remote partition of the second node into the first local partition of the first node in order to mirror changes to the first and second aggregates made while the first node had failed and the second node had taken over control of the first and second aggregates. Furthermore, the second node actively mirrors I/O data, directed to the first and second aggregates and being processed by the second node using the second remote partition, to the first local partition of the first node.

(18) In certain situations, data corruption can result during the transition of aggregates between nodes, such as during a giveback procedure or a planned takeover. Continuing with the example, if for some reason merely control of the first aggregate is transferred back to the first node and control of the second aggregate is retained by the second node, then data corruption/loss can occur. This is because the second remote partition of the second persistent memory of the second node is kept as an exact mirror of the first local partition of the first persistent memory of the first node. However, the first node is serving I/O directed to the first aggregate using the first local partition and the second node is serving I/O directed to the second aggregate using the second remote partition. It is possible that the first node, servicing I/O to the first aggregate using the first local partition, may write to a page block number (e.g., a free page within the first local partition corresponding to page block number **500**) that is the same page block number to which the second node is writing to while servicing I/O to the second aggregate using the second remote partition (e.g., a free page within the second remote partition corresponding to the page block number **500**). Data corruption occurs in this example because only one version of the data may be persisted (data loss) once the first local partition and the second remote partition are mirrored to one another so as to prevent a mismatch of data where page block number **500** of the first local partition has different data than page block number **500** of the second remote partition (data corruption).

(19) Accordingly, as provided herein, data corruption and data loss is avoided when transitioning control of aggregates between nodes during a giveback of aggregates between the nodes, a takeover (planned takeover) of aggregates between the nodes, or any other scenario where control of one or more aggregates is being transitioned between nodes. In an embodiment, when the second node determines that control of an aggregate (e.g., the first aggregate or the second aggregate originally owned/controlled by the first node) is to be given back from the second node to the first node, the second node allocates a portion of available free storage space within the second remote partition of the second persistent memory of the second node. The portion of available free storage space comprises a set of free pages within the second remote partition of the second persistent memory of the second node that are now being allocated/reserved for subsequent use when processing and/or mirroring data of I/O directed to the aggregate. The set of free pages are removed from a free pages list maintained by the second node for the second persistent memory, such as where page block numbers used to index the set of free pages are removed from the free pages list.

(20) The second node updates metadata information, corresponding to the second persistent memory, for the set of free pages with an identifier of the aggregate to create updated metadata information reserving the allocated pages for the aggregate. For example, the page block numbers used to index the set of free pages are tagged with the identifier of the aggregate so that the updated metadata information specifies that those page block numbers are being reserved for the aggregate. The second node transmits the updated metadata information to the first node.

(21) The first node uses the updated metadata information to build a free pages list of free pages for the aggregate to use from the first local partition. The first node builds the free pages list by traversing the updated metadata information and by adding entries tagged with the identifier of the aggregate to the free pages list. Any subsequent allocations of free pages from the first local partition to store I/O directed to the aggregate and processed by the first node will use free pages within the free pages list, corresponding to those pages that were previously reserved and available within the second remote partition of the second persistent memory of the second node. In an embodiment, the free pages list is used until all aggregates whose control are to be transitioned have successfully been transitioned from the second node to the first node. In this way, the first local partition and the second remote partition implement a temporary soft partition of free pages reserved for the aggregate during a transition window of transitioning control of the aggregates. The second node then transfers control of the aggregate to the first node.

(22) If for some reason control of the aggregate does not transfer to the first node, then both the first node and the second node have the same free pages lists to use for respectively reserving pages

within the first local partition and the second remote partition for subsequent allocation for the aggregate (e.g., the aggregate whose control was retained by the second node) and another aggregate (e.g., an aggregate whose control was passed to the first node) even though the second node will continue to process the I/O directed to the partition based upon retaining control of the partition. In this way, data corruption/loss is avoided because the same free pages (the same page block numbers) are reserved by both nodes for the partition. For example, control of a first aggregate may be retained by the second node (e.g., due to a failure to transition control of the first aggregate) but control of a second aggregate may be successfully transferred back to the first node. Without the reservation of the same free pages for storing data of the first aggregate, data corruption/loss could otherwise occur where the first node tries to allocate a free page for the second aggregate that is the same as a free page being allocated by the second node for the first aggregate.

(23) This partitioning of the first persistent memory and the second persistent memory may be performed during a transition window of transitioning control of aggregates between nodes. This partitioning may be referred to as a soft partition because it may be determined in real-time during the transition window as opposed to being a hard set partition. After the transition window, the soft partition may be removed such that there is normally no partition of storage between aggregates controlled by the same node (e.g., no partition within a local partition of persistent memory and/or no partition within a remote partition of the persistent memory), which provides improved flexibility for managing persistent memories of the nodes. Furthermore, the nodes do not have to maintain a centralized free list that is updated using inter-node communication, which simplifies the ability to avoid data corruption and loss, along with reducing inter-node communication for improving/reducing bandwidth overhead.

(24) FIG. 1 is a diagram illustrating an example operating environment **100** in which an embodiment of the techniques described herein may be implemented. In one example, the techniques described herein may be implemented within a client device **128**, such as a laptop, a tablet, a personal computer, a mobile device, a server, a virtual machine, a wearable device, etc. In another example, the techniques described herein may be implemented within one or more nodes, such as a first node **130** and/or a second node **132** within a first cluster **134**, a third node **136** within a second cluster **138**, etc. A node may comprise a storage controller, a server, an on-premise device, a virtual machine such as a storage virtual machine, hardware, software, or combination thereof. The one or more nodes may be configured to manage the storage and access to data on behalf of the client device **128** and/or other client devices. In another example, the techniques described herein may be implemented within a distributed computing platform **102** such as a cloud computing environment (e.g., a cloud storage environment, a multi-tenant platform, a hyperscale infrastructure comprising scalable server architectures and virtual networking, etc.) configured to manage the storage and access to data on behalf of client devices and/or nodes.

(25) In yet another example, at least some of the techniques described herein are implemented across one or more of the client device **128**, the one or more nodes **130**, **132**, and/or **136**, and/or the distributed computing platform **102**. For example, the client device **128** may transmit operations, such as data operations to read data and write data and metadata operations (e.g., a create file operation, a rename directory operation, a resize operation, a set attribute operation, etc.), over a network **126** to the first node **130** for implementation by the first node **130** upon storage. The first node **130** may store data associated with the operations within volumes or other data objects/structures hosted within locally attached storage, remote storage hosted by other computing devices accessible over the network **126**, storage provided by the distributed computing platform **102**, etc. The first node **130** may replicate the data and/or the operations to other computing devices, such as to the second node **132**, the third node **136**, a storage virtual machine executing within the distributed computing platform **102**, etc., so that one or more replicas of the data are maintained. For example, the third node **136** may host a destination storage volume that is

maintained as a replica of a source storage volume of the first node **130**. Such replicas can be used for disaster recovery and failover.

(26) In an embodiment, the techniques described herein are implemented by a storage operating system or are implemented by a separate module that interacts with the storage operating system. The storage operating system may be hosted by the client device, **128**, a node, the distributed computing platform **102**, or across a combination thereof. In an example, the storage operating system may execute within a storage virtual machine, a hyperscaler, or other computing environment. The storage operating system may implement a one or more file systems to logically organize data within storage devices as one or more storage objects and provide a logical/virtual representation of how the storage objects are organized on the storage devices (e.g., a file system tailored for block-addressable storage, a file system tailored for byte-addressable storage such as persistent memory). A storage object may comprise any logically definable storage element stored by the storage operating system (e.g., a volume stored by the first node **130**, a cloud object stored by the distributed computing platform **102**, etc.). Each storage object may be associated with a unique identifier that uniquely identifies the storage object. For example, a volume may be associated with a volume identifier uniquely identifying that volume from other volumes. The storage operating system also manages client access to the storage objects.

(27) The storage operating system may implement a file system for logically organizing data. For example, the storage operating system may implement a write anywhere file layout for a volume where modified data for a file may be written to any available location as opposed to a write-in-place architecture where modified data is written to the original location, thereby overwriting the previous data. In an example, the file system may be implemented through a file system layer that stores data of the storage objects in an on-disk format representation that is block-based (e.g., data is stored within 4 kilobyte blocks and inodes are used to identify files and file attributes such as creation time, access permissions, size and block location, etc.).

(28) In an example, deduplication may be implemented by a deduplication module associated with the storage operating system. Deduplication is performed to improve storage efficiency. One type of deduplication is inline deduplication that ensures blocks are deduplicated before being written to a storage device. Inline deduplication uses a data structure, such as an incore hash store, which maps fingerprints of data to data blocks of the storage device storing the data. Whenever data is to be written to the storage device, a fingerprint of that data is calculated and the data structure is looked up using the fingerprint to find duplicates (e.g., potentially duplicate data already stored within the storage device). If duplicate data is found, then the duplicate data is loaded from the storage device and a byte by byte comparison may be performed to ensure that the duplicate data is an actual duplicate of the data to be written to the storage device. If the data to be written is a duplicate of the loaded duplicate data, then the data to be written to disk is not redundantly stored to the storage device. Instead, a pointer or other reference is stored in the storage device in place of the data to be written to the storage device. The pointer points to the duplicate data already stored in the storage device. A reference count for the data may be incremented to indicate that the pointer now references the data. If at some point the pointer no longer references the data (e.g., the deduplicated data is deleted and thus no longer references the data in the storage device), then the reference count is decremented. In this way, inline deduplication is able to deduplicate data before the data is written to disk. This improves the storage efficiency of the storage device.

(29) Background deduplication is another type of deduplication that deduplicates data already written to a storage device. Various types of background deduplication may be implemented. In an example of background deduplication, data blocks that are duplicated between files are rearranged within storage units such that one copy of the data occupies physical storage. References to the single copy can be inserted into a file system structure such that all files or containers that contain the data refer to the same instance of the data. Deduplication can be performed on a data storage device block basis. In an example, data blocks on a storage device can be identified using a

physical volume block number. The physical volume block number uniquely identifies a particular block on the storage device. Additionally, blocks within a file can be identified by a file block number. The file block number is a logical block number that indicates the logical position of a block within a file relative to other blocks in the file. For example, file block number 0 represents the first block of a file, file block number 1 represents the second block, etc. File block numbers can be mapped to a physical volume block number that is the actual data block on the storage device. During deduplication operations, blocks in a file that contain the same data are deduplicated by mapping the file block number for the block to the same physical volume block number, and maintaining a reference count of the number of file block numbers that map to the physical volume block number. For example, assume that file block number 0 and file block number 5 of a file contain the same data, while file block numbers 1-4 contain unique data. File block numbers 1-4 are mapped to different physical volume block numbers. File block number 0 and file block number 5 may be mapped to the same physical volume block number, thereby reducing storage requirements for the file. Similarly, blocks in different files that contain the same data can be mapped to the same physical volume block number. For example, if file block number 0 of file A contains the same data as file block number 3 of file B, file block number 0 of file A may be mapped to the same physical volume block number as file block number 3 of file B.

(30) In another example of background deduplication, a changelog is utilized to track blocks that are written to the storage device. Background deduplication also maintains a fingerprint database (e.g., a flat metafile) that tracks all unique block data such as by tracking a fingerprint and other filesystem metadata associated with block data. Background deduplication can be periodically executed or triggered based upon an event such as when the changelog fills beyond a threshold. As part of background deduplication, data in both the changelog and the fingerprint database is sorted based upon fingerprints. This ensures that all duplicates are sorted next to each other. The duplicates are moved to a dup file. The unique changelog entries are moved to the fingerprint database, which will serve as duplicate data for a next deduplication operation. In order to optimize certain filesystem operations needed to deduplicate a block, duplicate records in the dup file are sorted in certain filesystem semantic order (e.g., inode number and block number). Next, the duplicate data is loaded from the storage device and a whole block byte by byte comparison is performed to make sure duplicate data is an actual duplicate of the data to be written to the storage device. After, the block in the changelog is modified to point directly to the duplicate data as opposed to redundantly storing data of the block.

(31) In an example, deduplication operations performed by a data deduplication layer of a node can be leveraged for use on another node during data replication operations. For example, the first node **130** may perform deduplication operations to provide for storage efficiency with respect to data stored on a storage volume. The benefit of the deduplication operations performed on first node **130** can be provided to the second node **132** with respect to the data on first node **130** that is replicated to the second node **132**. In some aspects, a data transfer protocol, referred to as the LRSE (Logical Replication for Storage Efficiency) protocol, can be used as part of replicating consistency group differences from the first node **130** to the second node **132**. In the LRSE protocol, the second node **132** maintains a history buffer that keeps track of data blocks that it has previously received. The history buffer tracks the physical volume block numbers and file block numbers associated with the data blocks that have been transferred from first node **130** to the second node **132**. A request can be made of the first node **130** to not transfer blocks that have already been transferred. Thus, the second node **132** can receive deduplicated data from the first node **130**, and will not need to perform deduplication operations on the deduplicated data replicated from first node **130**.

(32) In an example, the first node **130** may preserve deduplication of data that is transmitted from first node **130** to the distributed computing platform **102**. For example, the first node **130** may create an object comprising deduplicated data. The object is transmitted from the first node **130** to the distributed computing platform **102** for storage. In this way, the object within the distributed

computing platform **102** maintains the data in a deduplicated state. Furthermore, deduplication may be preserved when deduplicated data is transmitted/replicated/mirrored between the client device **128**, the first node **130**, the distributed computing platform **102**, and/or other nodes or devices.

(33) In an example, compression may be implemented by a compression module associated with the storage operating system. The compression module may utilize various types of compression techniques to replace longer sequences of data (e.g., frequently occurring and/or redundant sequences) with shorter sequences, such as by using Huffman coding, arithmetic coding, compression dictionaries, etc. For example, an uncompressed portion of a file may comprise “ggggnnnnnnqqqqqqqqqq”, which is compressed to become “4g6n10q”. In this way, the size of the file can be reduced to improve storage efficiency. Compression may be implemented for compression groups. A compression group may correspond to a compressed group of blocks. The compression group may be represented by virtual volume block numbers. The compression group may comprise contiguous or non-contiguous blocks.

(34) Compression may be preserved when compressed data is transmitted/replicated/mirrored between the client device **128**, a node, the distributed computing platform **102**, and/or other nodes or devices. For example, an object may be created by the first node **130** to comprise compressed data. The object is transmitted from the first node **130** to the distributed computing platform **102** for storage. In this way, the object within the distributed computing platform **102** maintains the data in a compressed state.

(35) In an example, various types of synchronization may be implemented by a synchronization module associated with the storage operating system. In an example, synchronous replication may be implemented, such as between the first node **130** and the second node **132**. It may be appreciated that the synchronization module may implement synchronous replication between any devices within the operating environment **100**, such as between the first node **130** of the first cluster **134** and the third node **136** of the second cluster **138** and/or between a node of a cluster and an instance of a node or virtual machine in the distributed computing platform **102**.

(36) As an example, during synchronous replication, the first node **130** may receive a write operation from the client device **128**. The write operation may target a file stored within a volume managed by the first node **130**. The first node **130** replicates the write operation to create a replicated write operation. The first node **130** locally implements the write operation upon the file within the volume. The first node **130** also transmits the replicated write operation to a synchronous replication target, such as the second node **132** that maintains a replica volume as a replica of the volume maintained by the first node **130**. The second node **132** will execute the replicated write operation upon the replica volume so that the file within the volume and the replica volume comprises the same data. After, the second node **132** will transmit a success message to the first node **130**. With synchronous replication, the first node **130** does not respond with a success message to the client device **128** for the write operation until both the write operation is executed upon the volume and the first node **130** receives the success message that the second node **132** executed the replicated write operation upon the replica volume.

(37) In another example, asynchronous replication may be implemented, such as between the first node **130** and the third node **136**. It may be appreciated that the synchronization module may implement asynchronous replication between any devices within the operating environment **100**, such as between the first node **130** of the first cluster **134** and the distributed computing platform **102**. In an example, the first node **130** may establish an asynchronous replication relationship with the third node **136**. The first node **130** may capture a baseline snapshot of a first volume as a point in time representation of the first volume. The first node **130** may utilize the baseline snapshot to perform a baseline transfer of the data within the first volume to the third node **136** in order to create a second volume within the third node **136** comprising data of the first volume as of the point in time at which the baseline snapshot was created.

(38) After the baseline transfer, the first node **130** may subsequently create snapshots of the first

volume over time. As part of asynchronous replication, an incremental transfer is performed between the first volume and the second volume. In particular, a snapshot of the first volume is created. The snapshot is compared with a prior snapshot that was previously used to perform the last asynchronous transfer (e.g., the baseline transfer or a prior incremental transfer) of data to identify a difference in data of the first volume between the snapshot and the prior snapshot (e.g., changes to the first volume since the last asynchronous transfer). Accordingly, the difference in data is incrementally transferred from the first volume to the second volume. In this way, the second volume will comprise the same data as the first volume as of the point in time when the snapshot was created for performing the incremental transfer. It may be appreciated that other types of replication may be implemented, such as semi-sync replication.

(39) In an embodiment, the first node **130** may store data or a portion thereof within storage hosted by the distributed computing platform **102** by transmitting the data within objects to the distributed computing platform **102**. In one example, the first node **130** may locally store frequently accessed data within locally attached storage. Less frequently accessed data may be transmitted to the distributed computing platform **102** for storage within a data storage tier **108**.

(40) The data storage tier **108** may store data within a service data store **120**, and may store client specific data within client data stores assigned to such clients such as a client (1) data store **122** used to store data of a client (1) and a client (N) data store **124** used to store data of a client (N). The data stores may be physical storage devices or may be defined as logical storage, such as a virtual volume, LUNs, or other logical organizations of data that can be defined across one or more physical storage devices. In another example, the first node **130** transmits and stores all client data to the distributed computing platform **102**. In yet another example, the client device **128** transmits and stores the data directly to the distributed computing platform **102** without the use of the first node **130**.

(41) The management of storage and access to data can be performed by one or more storage virtual machines (SVMs) or other storage applications that provide software as a service (SaaS) such as storage software services. In one example, an SVM may be hosted within the client device **128**, within the first node **130**, or within the distributed computing platform **102** such as by the application server tier **106**. In another example, one or more SVMs may be hosted across one or more of the client device **128**, the first node **130**, and the distributed computing platform **102**. The one or more SVMs may host instances of the storage operating system.

(42) In an example, the storage operating system may be implemented for the distributed computing platform **102**. The storage operating system may allow client devices to access data stored within the distributed computing platform **102** using various types of protocols, such as a Network File System (NFS) protocol, a Server Message Block (SMB) protocol and Common Internet File System (CIFS), and Internet Small Computer Systems Interface (iSCSI), and/or other protocols. The storage operating system may provide various storage services, such as disaster recovery (e.g., the ability to non-disruptively transition client devices from accessing a primary node that has failed to a secondary node that is taking over for the failed primary node), backup and archive function, replication such as asynchronous and/or synchronous replication, deduplication, compression, high availability storage, cloning functionality (e.g., the ability to clone a volume, such as a space efficient flex clone), snapshot functionality (e.g., the ability to create snapshots and restore data from snapshots), data tiering (e.g., migrating infrequently accessed data to slower/cheaper storage), encryption, managing storage across various platforms such as between on-premise storage systems and multiple cloud systems, etc.

(43) In one example of the distributed computing platform **102**, one or more SVMs may be hosted by the application server tier **106**. For example, a server (1) **116** is configured to host SVMs used to execute applications such as storage applications that manage the storage of data of the client (1) within the client (1) data store **122**. Thus, an SVM executing on the server (1) **116** may receive data and/or operations from the client device **128** and/or the first node **130** over the network **126**. The

SVM executes a storage application and/or an instance of the storage operating system to process the operations and/or store the data within the client (1) data store **122**. The SVM may transmit a response back to the client device **128** and/or the first node **130** over the network **126**, such as a success message or an error message. In this way, the application server tier **106** may host SVMs, services, and/or other storage applications using the server (1) **116**, the server (N) **118**, etc.

(44) A user interface tier **104** of the distributed computing platform **102** may provide the client device **128** and/or the first node **130** with access to user interfaces associated with the storage and access of data and/or other services provided by the distributed computing platform **102**. In an example, a service user interface **110** may be accessible from the distributed computing platform **102** for accessing services subscribed to by clients and/or nodes, such as data replication services, application hosting services, data security services, human resource services, warehouse tracking services, accounting services, etc. For example, client user interfaces may be provided to corresponding clients, such as a client (1) user interface **112**, a client (N) user interface **114**, etc. The client (1) can access various services and resources subscribed to by the client (1) through the client (1) user interface **112**, such as access to a web service, a development environment, a human resource application, a warehouse tracking application, and/or other services and resources provided by the application server tier **106**, which may use data stored within the data storage tier **108**.

(45) The client device **128** and/or the first node **130** may subscribe to certain types and amounts of services and resources provided by the distributed computing platform **102**. For example, the client device **128** may establish a subscription to have access to three virtual machines, a certain amount of storage, a certain type/amount of data redundancy, a certain type/amount of data security, certain service level agreements (SLAs) and service level objectives (SLOs), latency guarantees, bandwidth guarantees, access to execute or host certain applications, etc. Similarly, the first node **130** can establish a subscription to have access to certain services and resources of the distributed computing platform **102**.

(46) As shown, a variety of clients, such as the client device **128** and the first node **130**, incorporating and/or incorporated into a variety of computing devices may communicate with the distributed computing platform **102** through one or more networks, such as the network **126**. For example, a client may incorporate and/or be incorporated into a client application (e.g., software) implemented at least in part by one or more of the computing devices.

(47) Examples of suitable computing devices include personal computers, server computers, desktop computers, nodes, storage servers, nodes, laptop computers, notebook computers, tablet computers or personal digital assistants (PDAs), smart phones, cell phones, and consumer electronic devices incorporating one or more computing device components, such as one or more electronic processors, microprocessors, central processing units (CPU), or controllers. Examples of suitable networks include networks utilizing wired and/or wireless communication technologies and networks operating in accordance with any suitable networking and/or communication protocol (e.g., the Internet). In use cases involving the delivery of customer support services, the computing devices noted represent the endpoint of the customer support delivery process, i.e., the consumer's device.

(48) The distributed computing platform **102**, such as a multi-tenant business data processing platform or cloud computing environment, may include multiple processing tiers, including the user interface tier **104**, the application server tier **106**, and a data storage tier **108**. The user interface tier **104** may maintain multiple user interfaces, including graphical user interfaces and/or web-based interfaces. The user interfaces may include the service user interface **110** for a service to provide access to applications and data for a client (e.g., a "tenant") of the service, as well as one or more user interfaces that have been specialized/customized in accordance with user specific requirements (e.g., as discussed above), which may be accessed via one or more APIs.

(49) The service user interface **110** may include components enabling a tenant to administer the

tenant's participation in the functions and capabilities provided by the distributed computing platform **102**, such as accessing data, causing execution of specific data processing operations, etc. Each processing tier may be implemented with a set of computers, virtualized computing environments such as a storage virtual machine or storage virtual server, and/or computer components including computer servers and processors, and may perform various functions, methods, processes, or operations as determined by the execution of a software application or set of instructions.

(50) The data storage tier **108** may include one or more data stores, which may include the service data store **120** and one or more client data stores **122-124**. Each client data store may contain tenant-specific data that is used as part of providing a range of tenant-specific business and storage services or functions, including but not limited to ERP, CRM, eCommerce, Human Resources management, payroll, storage services, etc. Data stores may be implemented with any suitable data storage technology, including structured query language (SQL) based relational database management systems (RDBMS), file systems hosted by operating systems, object storage, etc.

(51) The distributed computing platform **102** may be a multi-tenant and service platform operated by an entity in order to provide multiple tenants with a set of business related applications, data storage, and functionality. These applications and functionality may include ones that a business uses to manage various aspects of its operations. For example, the applications and functionality may include providing web-based access to business information systems, thereby allowing a user with a browser and an Internet or intranet connection to view, enter, process, or modify certain types of business information or any other type of information.

(52) A clustered network environment **200** that may implement one or more aspects of the techniques described and illustrated herein is shown in FIG. 2. The clustered network environment **200** includes data storage apparatuses **202(1)-202(n)** that are coupled over a cluster or cluster fabric **204** that includes one or more communication network(s) and facilitates communication between the data storage apparatuses **202(1)-202(n)** (and one or more modules, components, etc. therein, such as, node computing devices **206(1)-206(n)**, for example), although any number of other elements or components can also be included in the clustered network environment **200** in other examples. This technology provides a number of advantages including methods, non-transitory computer readable media, and computing devices that implement the techniques described herein.

(53) In this example, node computing devices **206(1)-206(n)** can be primary or local storage controllers or secondary or remote storage controllers that provide client devices **208(1)-208(n)** with access to data stored within data storage devices **210(1)-210(n)** and cloud storage device(s) **236** (also referred to as cloud storage node(s)). The node computing devices **206(1)-206(n)** may be implemented as hardware, software (e.g., a storage virtual machine), or a combination thereof.

(54) The data storage apparatuses **202(1)-202(n)** and/or node computing devices **206(1)-206(n)** of the examples described and illustrated herein are not limited to any particular geographic areas and can be clustered locally and/or remotely via a cloud network, or not clustered in other examples. Thus, in one example the data storage apparatuses **202(1)-202(n)** and/or node computing device **206(1)-206(n)** can be distributed over a plurality of storage systems located in a plurality of geographic locations (e.g., located on-premise, located within a cloud computing environment, etc.); while in another example a clustered network can include data storage apparatuses **202(1)-202(n)** and/or node computing device **206(1)-206(n)** residing in a same geographic location (e.g., in a single on-site rack).

(55) In the illustrated example, one or more of the client devices **208(1)-208(n)**, which may be, for example, personal computers (PCs), computing devices used for storage (e.g., storage servers), or other computers or peripheral devices, are coupled to the respective data storage apparatuses **202(1)-202(n)** by network connections **212(1)-212(n)**. Network connections **212(1)-212(n)** may include a local area network (LAN) or wide area network (WAN) (i.e., a cloud network), for example, that utilize TCP/IP and/or one or more Network Attached Storage (NAS) protocols, such

as a Common Internet Filesystem (CIFS) protocol or a Network Filesystem (NFS) protocol to exchange data packets, a Storage Area Network (SAN) protocol, such as Small Computer System Interface (SCSI) or Fiber Channel Protocol (FCP), an object protocol, such as simple storage service (S3), and/or non-volatile memory express (NVMe), for example.

(56) Illustratively, the client devices **208(1)-208(n)** may be general-purpose computers running applications and may interact with the data storage apparatuses **202(1)-202(n)** using a client/server model for exchange of information. That is, the client devices **208(1)-208(n)** may request data from the data storage apparatuses **202(1)-202(n)** (e.g., data on one of the data storage devices **210(1)-210(n)** managed by a network storage controller configured to process I/O commands issued by the client devices **208(1)-208(n)**), and the data storage apparatuses **202(1)-202(n)** may return results of the request to the client devices **208(1)-208(n)** via the network connections **212(1)-212(n)**.

(57) The node computing devices **206(1)-206(n)** of the data storage apparatuses **202(1)-202(n)** can include network or host nodes that are interconnected as a cluster to provide data storage and management services, such as to an enterprise having remote locations, cloud storage (e.g., a storage endpoint may be stored within cloud storage device(s) **236**), etc., for example. Such node computing devices **206(1)-206(n)** can be attached to the cluster fabric **204** at a connection point, redistribution point, or communication endpoint, for example. One or more of the node computing devices **206(1)-206(n)** may be capable of sending, receiving, and/or forwarding information over a network communications channel, and could comprise any type of device that meets any or all of these criteria.

(58) In an example, the node computing devices **206(1)** and **206(n)** may be configured according to a disaster recovery configuration whereby a surviving node provides switchover access to the storage devices **210(1)-210(n)** in the event a disaster occurs at a disaster storage site (e.g., the node computing device **206(1)** provides client device **212(n)** with switchover data access to data storage devices **210(n)** in the event a disaster occurs at the second storage site). In other examples, the node computing device **206(n)** can be configured according to an archival configuration and/or the node computing devices **206(1)-206(n)** can be configured based on another type of replication arrangement (e.g., to facilitate load sharing). Additionally, while two node computing devices are illustrated in FIG. 2, any number of node computing devices or data storage apparatuses can be included in other examples in other types of configurations or arrangements. In an example, control of aggregates may be switched between the node computing devices **206(1)** and **206(n)** in the event of a disaster or planned takeover. As provided herein, block allocation for persistent memory during aggregate transition between the node computing devices **206(1)** and **206(n)** is performed.

(59) As illustrated in the clustered network environment **200**, node computing devices **206(1)-206(n)** can include various functional components that coordinate to provide a distributed storage architecture. For example, the node computing devices **206(1)-206(n)** can include network modules **214(1)-214(n)** and disk modules **216(1)-216(n)**. Network modules **214(1)-214(n)** can be configured to allow the node computing devices **206(1)-206(n)** (e.g., network storage controllers) to connect with client devices **208(1)-208(n)** over the storage network connections **212(1)-212(n)**, for example, allowing the client devices **208(1)-208(n)** to access data stored in the clustered network environment **200**.

(60) Further, the network modules **214(1)-214(n)** can provide connections with one or more other components through the cluster fabric **204**. For example, the network module **214(1)** of node computing device **206(1)** can access the data storage device **210(n)** by sending a request via the cluster fabric **204** through the disk module **216(n)** of node computing device **206(n)** when the node computing device **206(n)** is available. Alternatively, when the node computing device **206(n)** fails, the network module **214(1)** of node computing device **206(1)** can access the data storage device **210(n)** directly via the cluster fabric **204**. The cluster fabric **204** can include one or more local and/or wide area computing networks (i.e., cloud networks) embodied as Infiniband, Fibre Channel

(FC), or Ethernet networks, for example, although other types of networks supporting other protocols can also be used.

(61) Disk modules **216(1)-216(n)** can be configured to connect data storage devices **210(1)-210(n)**, such as disks or arrays of disks, SSDs, flash memory, or some other form of data storage, to the node computing devices **206(1)-206(n)**. Often, disk modules **216(1)-216(n)** communicate with the data storage devices **210(1)-210(n)** according to the SAN protocol, such as SCSI or FCP, for example, although other protocols can also be used. Thus, as seen from an operating system on node computing devices **206(1)-206(n)**, the data storage devices **210(1)-210(n)** can appear as locally attached. In this manner, different node computing devices **206(1)-206(n)**, etc. may access data blocks, files, or objects through the operating system, rather than expressly requesting abstract files.

(62) While the clustered network environment **200** illustrates an equal number of network modules **214(1)-214(n)** and disk modules **216(1)-216(n)**, other examples may include a differing number of these modules. For example, there may be a plurality of network and disk modules interconnected in a cluster that do not have a one-to-one correspondence between the network and disk modules. That is, different node computing devices can have a different number of network and disk modules, and the same node computing device can have a different number of network modules than disk modules.

(63) Further, one or more of the client devices **208(1)-208(n)** can be networked with the node computing devices **206(1)-206(n)** in the cluster, over the storage connections **212(1)-212(n)**. As an example, respective client devices **208(1)-208(n)** that are networked to a cluster may request services (e.g., exchanging of information in the form of data packets) of node computing devices **206(1)-206(n)** in the cluster, and the node computing devices **206(1)-206(n)** can return results of the requested services to the client devices **208(1)-208(n)**. In one example, the client devices **208(1)-208(n)** can exchange information with the network modules **214(1)-214(n)** residing in the node computing devices **206(1)-206(n)** (e.g., network hosts) in the data storage apparatuses **202(1)-202(n)**.

(64) In one example, the storage apparatuses **202(1)-202(n)** host aggregates corresponding to physical local and remote data storage devices, such as local flash or disk storage in the data storage devices **210(1)-210(n)**, for example. One or more of the data storage devices **210(1)-210(n)** can include mass storage devices, such as disks of a disk array. The disks may comprise any type of mass storage devices, including but not limited to magnetic disk drives, flash memory, and any other similar media adapted to store information, including, for example, data and/or parity information. In an example, control of the aggregates stored within the data storage devices **210(1)-210(n)** may be switched between the node computing devices **206(1)** and **206(n)** in the event of a disaster or planned takeover. As provided herein, block allocation for persistent memory during aggregate transition between the node computing devices **206(1)** and **206(n)** is performed.

(65) The aggregates include volumes **218(1)-218(n)** in this example, although any number of volumes can be included in the aggregates. The volumes **218(1)-218(n)** are virtual data stores or storage objects that define an arrangement of storage and one or more filesystems within the clustered network environment **200**. Volumes **218(1)-218(n)** can span a portion of a disk or other storage device, a collection of disks, or portions of disks, for example, and typically define an overall logical arrangement of data storage. In one example volumes **218(1)-218(n)** can include stored user data as one or more files, blocks, or objects that may reside in a hierarchical directory structure within the volumes **218(1)-218(n)**.

(66) Volumes **218(1)-218(n)** are typically configured in formats that may be associated with particular storage systems, and respective volume formats typically comprise features that provide functionality to the volumes **218(1)-218(n)**, such as providing the ability for volumes **218(1)-218(n)** to form clusters, among other functionality. Optionally, one or more of the volumes **218(1)-218(n)** can be in composite aggregates and can extend between one or more of the data

storage devices **210(1)-210(n)** and one or more of the cloud storage device(s) **236** to provide tiered storage, for example, and other arrangements can also be used in other examples.

(67) In one example, to facilitate access to data stored on the disks or other structures of the data storage devices **210(1)-210(n)**, a filesystem may be implemented that logically organizes the information as a hierarchical structure of directories and files. In this example, respective files may be implemented as a set of disk blocks of a particular size that are configured to store information, whereas directories may be implemented as specially formatted files in which information about other files and directories are stored.

(68) Data can be stored as files or objects within a physical volume and/or a virtual volume, which can be associated with respective volume identifiers. The physical volumes correspond to at least a portion of physical storage devices, such as the data storage devices **210(1)-210(n)** (e.g., a Redundant Array of Independent (or Inexpensive) Disks (RAID system)) whose address, addressable space, location, etc. does not change. Typically the location of the physical volumes does not change in that the range of addresses used to access it generally remains constant.

(69) Virtual volumes, in contrast, can be stored over an aggregate of disparate portions of different physical storage devices. Virtual volumes may be a collection of different available portions of different physical storage device locations, such as some available space from disks, for example. It will be appreciated that since the virtual volumes are not “tied” to any one particular storage device, virtual volumes can be said to include a layer of abstraction or virtualization, which allows it to be resized and/or flexible in some regards.

(70) Further, virtual volumes can include one or more logical unit numbers (LUNs), directories, Qtrees, files, and/or other storage objects, for example. Among other things, these features, but more particularly the LUNs, allow the disparate memory locations within which data is stored to be identified, for example, and grouped as data storage unit. As such, the LUNs may be characterized as constituting a virtual disk or drive upon which data within the virtual volumes is stored within an aggregate. For example, LUNs are often referred to as virtual drives, such that they emulate a hard drive, while they actually comprise data blocks stored in various parts of a volume.

(71) In one example, the data storage devices **210(1)-210(n)** can have one or more physical ports, wherein each physical port can be assigned a target address (e.g., SCSI target address). To represent respective volumes, a target address on the data storage devices **210(1)-210(n)** can be used to identify one or more of the LUNs. Thus, for example, when one of the node computing devices **206(1)-206(n)** connects to a volume, a connection between the one of the node computing devices **206(1)-206(n)** and one or more of the LUNs underlying the volume is created.

(72) Respective target addresses can identify multiple of the LUNs, such that a target address can represent multiple volumes. The I/O interface, which can be implemented as circuitry and/or software in a storage adapter or as executable code residing in memory and executed by a processor, for example, can connect to volumes by using one or more addresses that identify the one or more of the LUNs.

(73) Referring to FIG. 3, node computing device **206(1)** in this particular example includes processor(s) **300**, a memory **302**, a network adapter **304**, a cluster access adapter **306**, and a storage adapter **308** interconnected by a system bus **310**. In other examples, the node computing device **206(1)** comprises a virtual machine, such as a virtual storage machine. The node computing device **206(1)** also includes a storage operating system **312** installed in the memory **302** that can, for example, implement a RAID data loss protection and recovery scheme to optimize reconstruction of data of a failed disk or drive in an array, along with other functionality such as deduplication, compression, snapshot creation, data mirroring, synchronous replication, asynchronous replication, encryption, etc. In some examples, the node computing device **206(n)** is substantially the same in structure and/or operation as node computing device **206(1)**, although the node computing device **206(n)** can also include a different structure and/or operation in one or more aspects than the node computing device **206(1)**. In an example, a file system may be implemented for persistent memory.

- (74) The network adapter **304** in this example includes the mechanical, electrical and signaling circuitry needed to connect the node computing device **206(1)** to one or more of the client devices **208(1)-208(n)** over network connections **212(1)-212(n)**, which may comprise, among other things, a point-to-point connection or a shared medium, such as a local area network. In some examples, the network adapter **304** further communicates (e.g., using TCP/IP) via the cluster fabric **204** and/or another network (e.g. a WAN) (not shown) with cloud storage device(s) **236** to process storage operations associated with data stored thereon.
- (75) The storage adapter **308** cooperates with the storage operating system **312** executing on the node computing device **206(1)** to access information requested by one of the client devices **208(1)-208(n)** (e.g., to access data on a data storage device **210(1)-210(n)** managed by a network storage controller). The information may be stored on any type of attached array of writeable media such as magnetic disk drives, flash memory, and/or any other similar media adapted to store information.
- (76) In the exemplary data storage devices **210(1)-210(n)**, information can be stored in data blocks on disks. The storage adapter **308** can include I/O interface circuitry that couples to the disks over an I/O interconnect arrangement, such as a storage area network (SAN) protocol (e.g., Small Computer System Interface (SCSI), Internet SCSI (iSCSI), hyperSCSI, Fiber Channel Protocol (FCP)). The information is retrieved by the storage adapter **308** and, if necessary, processed by the processor(s) **300** (or the storage adapter **308** itself) prior to being forwarded over the system bus **310** to the network adapter **304** (and/or the cluster access adapter **306** if sending to another node computing device in the cluster) where the information is formatted into a data packet and returned to a requesting one of the client devices **208(1)-208(n)** and/or sent to another node computing device attached via the cluster fabric **204**. In some examples, a storage driver **314** in the memory **302** interfaces with the storage adapter to facilitate interactions with the data storage devices **210(1)-210(n)**.
- (77) The storage operating system **312** can also manage communications for the node computing device **206(1)** among other devices that may be in a clustered network, such as attached to a cluster fabric **204**. Thus, the node computing device **206(1)** can respond to client device requests to manage data on one of the data storage devices **210(1)-210(n)** or cloud storage device(s) **236** (e.g., or additional clustered devices) in accordance with the client device requests.
- (78) The file system module **318** of the storage operating system **312** can establish and manage one or more filesystems including software code and data structures that implement a persistent hierarchical namespace of files and directories, for example. As an example, when a new data storage device (not shown) is added to a clustered network system, the file system module **318** is informed where, in an existing directory tree, new files associated with the new data storage device are to be stored. This is often referred to as “mounting” a filesystem.
- (79) In the example node computing device **206(1)**, memory **302** can include storage locations that are addressable by the processor(s) **300** and adapters **304**, **306**, and **308** for storing related software application code and data structures. The processor(s) **300** and adapters **304**, **306**, and **308** may, for example, include processing elements and/or logic circuitry configured to execute the software code and manipulate the data structures.
- (80) In the example, the node computing device **206(1)** comprises persistent memory **320**. The persistent memory **320** comprises a plurality of pages within which data can be stored. The plurality of pages may be indexed by page block numbers.
- (81) The storage operating system **312**, portions of which are typically resident in the memory **302** and executed by the processor(s) **300**, invokes storage operations in support of a file service implemented by the node computing device **206(1)**. Other processing and memory mechanisms, including various computer readable media, may be used for storing and/or executing application instructions pertaining to the techniques described and illustrated herein. For example, the storage operating system **312** can also utilize one or more control files (not shown) to aid in the

provisioning of virtual machines.

(82) In this particular example, the memory **302** also includes a module configured to implement the techniques described herein, including for example block allocation for persistent memory during aggregate transition as discussed above and further below.

(83) The examples of the technology described and illustrated herein may be embodied as one or more non-transitory computer or machine readable media, such as the memory **302**, having machine or processor-executable instructions stored thereon for one or more aspects of the present technology, which when executed by processor(s), such as processor(s) **300**, cause the processor(s) to carry out the steps necessary to implement the methods of this technology, as described and illustrated with the examples herein. In some examples, the executable instructions are configured to perform one or more steps of a method described and illustrated later.

(84) One embodiment of block allocation for persistent memory during aggregate transition is illustrated by an exemplary method **400** of FIG. **4**, which is further described in conjunction with system **500** of FIGS. **5A-5G**. A node (A) **502** may be paired with a node (B) **524**, such as where the node (A) **502** and the node (B) **524** are a high availability node pair that can provide takeover and giveback functionality in order to provide clients with non-disruptive access to data in the event one of the nodes fails. The node (A) **502** and/or the node (B) **524** may be implemented as a server, a virtual machine (e.g., a storage virtual machine hosted within a cloud computing environment), software such as software as a service (SaaS), hardware, or a combination thereof.

(85) The node (A) **502** may originally host and control an aggregate (A1) **518**, an aggregate (A2) **520**, and/or other local aggregates within which the node (A) **502** may store data on behalf of client devices, as illustrated by FIG. **5A**. The node (A) **502** may store at least some data of the aggregate (A1) **518** within a first local partition **506** of a persistent memory (A) **504** of the node (A) **502** as aggregate (A1) data **510**. For example, the persistent memory (A) **504** may provide lower latency and improved performance compared to other storage available to the node (A) **502**, such as disk drives, cloud storage, solid state drives, etc. Thus, certain data may be stored within the persistent memory (A) **504** such as recently accessed data, frequently accessed data, data predicted to be accessed within a threshold time span, etc. The node (A) **502** may store at least some data of the aggregate (A2) **520** within the first local partition **506** of the persistent memory (A) **504** as aggregate (A2) data **512**. During normal operation, the first local partition **506** may not be partitioned, such that there is no hard partition to separate/divide the first local partition **506** for separately storing and partitioning the aggregate (A1) data **510** from the aggregate (A2) data **512**.

(86) The node (B) **524** may originally host and control an aggregate (B1) **542**, an aggregate (B2) **540**, and/or other local aggregates within which the node (B) **524** may store data on behalf of client devices. The node (B) **524** may store at least some data of the aggregate (B1) **542** within a second local partition **528** of a persistent memory (B) **526** of the node (B) **524** as aggregate (B1) data **532**. For example, the persistent memory (B) **526** may provide lower latency and improved performance compared to other storage available to the node (B) **524**, such as disk drives, cloud storage, solid state drives, etc. Thus, certain data may be stored within the persistent memory (B) **526** such as recently accessed data, frequently accessed data, data predicted to be accessed within a threshold time span, etc. The node (B) **524** may store at least some data of the aggregate (B2) **540** within the second local partition **528** of the persistent memory (B) **526** as aggregate (B2) data **534**. During normal operation, the second local partition **528** may not be partitioned, such that there is no hard partition to separate/divide the second local partition **528** for separately storing and partitioning the aggregate (B1) data **532** from the aggregate (B2) data **534**.

(87) In order to provide client with non-disruptive access to data within the aggregates hosted by the nodes, the nodes need access to up-to-date data of the other nodes, such as data stored within a persistent memory of a partner node. Accordingly, the persistent memory (A) **504** of the node (A) **502** comprises a first remote partition **508**. The aggregate (B1) data **532** is mirrored from the second local partition **528** of the persistent memory (B) **526** of the node (B) **524** into the first

remote partition **508** of the persistent memory (A) **504** of the node (A) **502** as mirrored aggregate (B1) data **514**. The aggregate (B2) data **534** is mirrored from the second local partition **528** of the persistent memory (B) **526** of the node (B) **524** into the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502** as mirrored aggregate (B2) data **516**.

(88) Similarly, the persistent memory (B) **526** of the node (B) **524** comprises a second remote partition **530**. The aggregate (A1) data **510** is mirrored from the first local partition **506** of the persistent memory (A) **504** of the node (A) **502** into the second remote partition **530** of the persistent memory (B) **526** of the node (B) **524** as mirrored aggregate (A1) data **536**. The aggregate (A2) data **512** is mirrored from the first local partition **506** of the persistent memory (A) **504** of the node (A) **502** into the second remote partition **530** of the persistent memory (B) **526** of the node (B) **524** as mirrored aggregate (A2) data **538**.

(89) The node (A) **502** may maintain various metadata information **522** regarding the persistent memory (A) **504**. In an embodiment, the metadata information **522** may indicate whether a page within the persistent memory (A) **504** is reserved for use by an aggregate. In another embodiment, the node (A) **502** may maintain free lists of free pages within the first local partition **506** and the first remote partition **508** that are available to allocate for storing data of the aggregates (e.g., a list of page block numbers indexing free pages that do not comprise data referenced by at least one of an active file system or snapshots of the active file system).

(90) As the node (A) **502** processes I/O directed to the aggregate (A1) **518** and/or the aggregate (A2) **520** using the first local partition **506**, data associated with the I/O (e.g., data being written to the aggregate (A1) **518** and/or the aggregate (A2) **520**) is mirrored into the second remote partition **530** of the persistent memory (B) **526** of the node (B) **524**. Thus, if the node (A) **502** fails, the node (B) **524** will have access to up-to-date data of the aggregate (A1) **518** and the aggregate (A2) **520** within the second remote partition **530** for taking over subsequent processing I/O directed to the aggregate (A1) **518** and the aggregate (A2) **520** in place of the failed node (A) **502**.

(91) Similarly, as the node (B) **524** processes I/O directed to the aggregate (B1) **542** and/or the aggregate (B2) **540** using the second local partition **532**, data associated with the I/O (e.g., data being written to the aggregate (B1) **542** and/or the aggregate (B2) **540**) is mirrored into the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502**. Thus, if the node (B) **524** fails, the node (A) **502** will have access to up-to-date data of the aggregate (B1) **542** and the aggregate (B2) **540** within the first remote partition **508** for taking over subsequent processing I/O directed to the aggregate (B1) **542** and the aggregate (B2) **540** in place of the failed node (B) **524**.

(92) FIG. 5B illustrates an embodiment of the node (B) **524** failing **550**. In an embodiment, the node (A) **502** may detect the failure **550** of the node (B) **524** through various mechanisms, such as by detecting a loss of a heartbeat signal otherwise generated by the node (B) **524** during normal operation of the node (B) **524**. It may be appreciated that various mechanism may be used to determine whether a node has failed or is operational. In response to detecting the failure **550** of the node (B) **524**, the node (A) **502** performs a takeover procedure to take over control of the aggregate (B1) **542** and/or the aggregate (B2) **540** from the node (B) **524**, as illustrated by FIG. 5C. Because the node (B) **524** has failed, the node (A) **502** will take control of the aggregate (B1) **542** and the aggregate (B2) **540** so that the node (A) **502** can process I/O directed to the aggregate (B1) **542** and the aggregate (B2) **540** that would otherwise have been processed by the node (B) **524** during normal operation. As part of the takeover procedure, the node (A) **502** evaluates the mirrored aggregate (B1) data **514** and the mirrored aggregate (B2) data **516** within the first remote partition **508** of the persistent memory (A) **504** to build a free pages list of free pages (e.g., a list of page block numbers of the free pages) of the first remote partition **508** that are available for subsequent allocation by the node (A) **502** for storing data of I/O directed to the aggregate (B1) **542** and/or the aggregate (B2) **540**.

(93) Once the takeover procedure has completed, the node (A) **502** processes I/O directed to the aggregate (A1) **518** using the first local partition **506** of the persistent memory (A) **504** of the node

(A) 502. The node (A) 502 processes I/O directed to the aggregate (A2) 520 using the first local partition 506 of the persistent memory (A) 504 of the node (A) 502. The node (A) 502 processes I/O directed to the aggregate (B1) 542 using the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. The node (A) 502 processes I/O directed to the aggregate (B2) 540 using the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. In this way, the aggregate (A1) data 510, the aggregate (A2) data 512, the aggregate (B1) data 514, and/or the aggregate (B2) data 516 may change over time based upon the node (A) 502 processing write commands, delete commands, and/or other commands that modify such data.

(94) At some point in time, the node (B) 524 recovers from the failure 550, as illustrated by FIG. 5D. At 402 (of FIG. 4's exemplary method 400), a determination is made that the node (A) 502 is to transition control of the aggregate (B1) 542 and/or the aggregate (B2) 540 back from the node (A) 502 to the node (B) 524. While embodiments of FIG. 4's method 400 are described with respect to a giveback procedure, they may also apply to takeover procedures. However, discussion of FIG. 4's method 400 will be done with respect to a giveback procedure for sake of demonstration. Various indicators may trigger the determination that the node (A) 502 is to transition control of the aggregate (B1) 542 and/or the aggregate (B2) 540 to the node (B) 524. In an embodiment, the node (B) 524 may transmit a request to the node (A) 502 to perform a giveback procedure of the aggregate (B1) 542 and/or the aggregate (B2) 540. In another embodiment, the node (A) 502 may detect that the node (B) 524 has recovered from the failure 550, and thus inquire with node (B) 524 as to whether a giveback procedure of the aggregate (B1) 542 and/or the aggregate (B2) 540 should be performed. In this way, the node (A) 502 may determine that the giveback procedure is to be implemented to return control of the aggregate (B1) 542 and/or the aggregate (B2) 540 from the node (A) 502 to the node (B) 524.

(95) Once the node (B) 524 has recovered from the failure 550 (e.g., before the giveback procedure has commenced or completed), current data of the mirrored aggregate (B1) data 514 and/or the mirrored aggregate (B2) data 516 is resynchronized from the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502 into the second local partition 528 of the persistent memory (B) 526 of the node (B) 524. In an embodiment, a byte-by-byte resynchronization is performed to resynchronize changes made to the aggregate (B1) 542 and/or the aggregate (B2) 540 while controlled by the node (A) 502 to the second local partition 528 of the persistent memory (B) 526 of the node (B) 524. In this way, a resynchronization process is performed to resynchronize current data of the mirrored aggregate (B1) data 514 and/or the mirrored aggregate (B2) data 516 from the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502 into the second local partition 528 of the persistent memory (B) 526 of the node (B) 524.

(96) Furthermore, mirroring of incoming I/O operations directed to the aggregate (B1) 541 and/or the aggregate (B2) 540 is performed (e.g., based on the node (A) 502 knowing that node (B) 524 is recovered). In this way, the mirrored aggregate (B1) data 514 of the aggregate (B1) 542 and/or the mirrored aggregate (B2) data 516 of the aggregate (B2) 540 being modified by I/O operations being processed by the node (A) 502 using the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502 is being mirrored to the second local partition 528 of the persistent memory (B) 526 of the node (B) 524. For example, the node (A) 502 may receive a write operation from a client device before the giveback procedure has completed. The write operation may target the aggregate (B1) 542 in order to write data to the aggregate (B1) 542. Because the node (A) 502 has control over the aggregate (B1) 542, the node (A) 502 writes the data to the mirrored aggregate (B1) data 514 within the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. The node (A) 502 also mirrors that data into the aggregate (B1) data 532 within the second local partition 528 of the persistent memory (B) 526 of the node (B) 524 as part of processing the write operation. In this way, data of incoming operations is mirrored to the second local partition 528 of the persistent memory (B) 526 of the node (B) 524 (e.g., before the giveback procedure has commenced or completed).

(97) Once the node (A) 502 has determined that the giveback procedure is to be implemented to give back control of the aggregate (B1) 542 and/or the aggregate (B2) 540 from the node (A) 502 to the node (B) 524, the node (A) 502 may initiate giveback of the aggregate (B1) 542 and/or the aggregate (B2) 540 one at a time. In an embodiment, the node (A) 502 initiates giveback of the aggregate (B1) 542 to the node (B) 524. As part of the giveback, soft partitions (e.g., the allocation/reservation of certain free pages for use by the aggregate (B1) 542 and the allocation/reservation of different free pages for use by the aggregate (B2) 540) are created in the event control of one of the aggregates does not get transitioned to the node (B) 524, and thus the node (B) 524 would be serving I/O for that aggregate while the node (A) 502 would be serving I/O for the other aggregate. In that event, each node will now allocate different free pages to store data of the aggregate that node controls, thus avoiding instances of data corruption and loss where each node could otherwise allocate the same free page to store different data of the different aggregates.

(98) At 404, the node (A) 502 allocates a portion of available free storage space from the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. For example, the persistent memory (A) 504 may be comprised of pages within which data can be stored. The pages may be indexed by page block numbers (e.g., a first page having a first page block number, a second page having a second page block number, etc.). A page may be a free page that is available for storing data, such as because the free page does not comprise data or comprises data that is no longer referenced by an active file system and snapshots of the active file system. A page may be a used page that is unavailable for storing data, such as because the used page comprises data currently referenced by an active file system and/or one or more snapshots of the active file system.

(99) The portion of available free storage space may be allocated as a set of free pages within the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. The set of free pages may comprise free pages that reserved for subsequent allocation and use for storing data associated with the aggregate (B1) 542 within the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502. The node (A) 502 may determine what percentage of free pages within the first remote partition 508 to allocate and reserve as the portion of available free storage space for subsequent use to store data of the aggregate (B1) 542 based upon various factors such as how many other aggregates are using the first remote partition 508 for storing data. In an example, free storage space may be allocated evenly across all aggregates using the first remote partition 508. In another example, the percentage of free pages allocated for the aggregate (B1) 542 may be based upon historic storage space utilization by the aggregate (B1) 542 (e.g., if only a small percentage is historically used by the aggregate (B1) 542, then a relatively small percentage of the available free pages may be allocated/reserved for use by the aggregate (B1) 542), predicted utilization, etc. The set of free pages allocated/reserved for use by the aggregate (B1) 542 are removed from a free pages list maintained by the node (A) 502 of free pages available for use by other aggregates.

(100) In an embodiment, because the set of free pages is merely allocated/reserved during the transition of control of aggregates from the node (A) 502 to the node (B) 524, the first remote partition 508 and the second local partition 528 have a temporary soft partition of the set of free pages for the aggregate (B1) 540. In an embodiment, this soft partition can be removed once control of the aggregate(s) has been fully transferred and onlined (e.g., made available for access by client devices) by the node (B) 524.

(101) Implementing soft partitions, instead of hard partitions, for data of aggregates using the local partitions and remote partitions of the persistent memory (A) 504 of the node (A) 502 and the persistent memory (B) 526 of the node (B) 524 allows for the local partitions and the remote partitions to be sized and resized based upon various considerations (e.g., arbitrarily sized, sized/resized based upon current/historic/predicted utilization, etc.). Furthermore, implementing soft partitions, instead of hard partitions, for data of aggregates using the local partitions and remote partitions of the persistent memory (A) 504 of the node (A) 502 and the persistent memory

(B) 526 of the node (B) 524 allows for any number of aggregates to be supported. Thus, a node may store data of any number of aggregates within a local partition and/or a remote partition of persistent memory of the node. In an embodiment, because a soft partition is merely implemented for the local partitions and remote partitions of the persistent memories during a transition window of transitioning control of aggregates between nodes, no soft partition may be used outside of the transition window (e.g., the first local partition 506 is not partitioned for use by certain aggregates, the first remote partition 508 is not partitioned for use by certain aggregates, the second local partition 528 is not partitioned for use by certain aggregates, and the second remote partition 530 is not partitioned for use by certain aggregates when aggregates are not being transitioned between nodes). In an embodiment, once all aggregates have been transitioned, then any soft partitions may be removed.

(102) Once the portion of available free storage space of the first remote partition 508 of the persistent memory (A) 504 of the node (A) 502 has been allocated/reserved as allocated pages for the aggregate (B1) 542, the node (A) 502 may update the metadata information 522 as updated metadata information 552 for the aggregate (B1) 542 to indicate that the allocated pages are allocated/reserved for use in storing data of the aggregate (B1) 542, at 406. For example, the allocated pages that are allocated/reserved for the aggregate (B1) 542 are tagged with an identifier of the aggregate (B1) 542 to reserve the allocated pages for the aggregate (B1) 542 (e.g., page block numbers of the allocated pages may be tagged with the identifier). In this way, the updated metadata information 552 for the aggregate (B1) 542 comprises page block numbers of the allocated pages that are tagged with the identifier of the aggregate (B1) 542.

(103) At 408, the node (A) 502 mirrors the updated metadata information 552 for the aggregate (B1) 542 to the node (B) 524. For example, the updated metadata information 552 is transmitted from the node (A) 502 to the node (B) 524, as illustrated by FIG. 5D. In this way, the node (B) 524 may store the updated metadata information 552 for the aggregate (B1) 542 so that the node (B) 524 can identify what free pages within the second local partition 528 of the persistent memory (B) 526 of the node (B) 524 to reserve for use by the aggregate (B1) 542 once control of the aggregate (B1) 542 has been transitioned from the node (A) 502 to the node (B) 524. In this way, the node (B) 524 may construct a free pages list for the aggregate (B1) 542 (e.g., a list of page block numbers of free pages that are allocated/reserved for use by the node (B) 524 for storing data of the aggregate (B1) 542) by adding free pages (e.g., page block numbers of the free pages) tagged with the identifier of the aggregate (B1) 542 into the free pages list for the aggregate (B1) 542 and excluding other pages from the free pages list for the aggregate (B1) 542. Thus, the allocated pages, tagged with the identifier of the aggregate (B1) 542 within the updated metadata 522 for the aggregate (B1) 542, are added into the free pages list as being allocated/reserved for use to store data of the aggregate (B1) 542 within the second local partition 528 of the persistent memory (B) 526 of the node (B) 524. In this way, both the node (A) 502 and the node (B) 524 have reserved the same free pages for use in storing data of the aggregate (B1) 542.

(104) New block allocations for the aggregate (B1) 542 (e.g., allocations of free pages to store data of the aggregate (B1) 542) are directed to the allocated pages within the free pages list that were added to the free pages list based upon the allocated pages being tagged with the identifier of the aggregate (B1) 542 (e.g., until completion of the giveback of all aggregates). Because both the node (A) 502 and the node (B) 524 have the same allocated pages that are reserved for use by the aggregate (B1) 542 to store data of the aggregate (B1) 542 into the second local partition 528 of the persistent memory (B) 526 of the node (B) 524 and into the first remote partition 508 of the persistent memory (A) 508 of the node (A) 502, data corruption and loss is avoided. Data corruption is avoided where the node (A) 502 and the node (B) 524 could otherwise allocate the same page having the same page block number to store different data (e.g., node (A) 502 could store data of the aggregate (B2) 540 into a free page having a same page block number as a free page at which node (B) 524 stores data of the aggregate (B1) 524 after control of the aggregate

(B1) **542** has been transitioned to the node (B) **524**) if the node (A) **502** and the node (B) **524** did not have the same allocated pages for the aggregate (B1) **542**. Data corruption would result because the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502** and the second local partition **528** of the persistent memory (B) **526** of the node (B) **524** are to be mirrors of one another comprising the exact same data. Either the free page within the first remote partition **508** and the free page within the second local partition **528** will comprise different data, or due to mirroring one of the free pages will be overwritten and thus result in data loss.

(105) However, since both the node (A) **502** and the node (B) **524** have the same allocated pages for the aggregate (B1) **542**, data corruption and loss is avoided.

(106) At **410**, control of the aggregate (B1) **542** is given back **554** from the node (A) **502** to the node (B) **524**, as illustrated by FIG. 5E. In an embodiment, control of the aggregate (B1) **542** is given back **554** to the node (B) **524** upon confirmation that the node (B) **524** has constructed the free pages list for the aggregate (B1) **542** based upon the updated metadata information **552** for the aggregate (B1) **542**. Once node (B) **524** has control of the aggregate (B1) **542**, the node (B) **524** will actively serve I/O directed to the aggregate (B1) **542** using the allocated pages, within the free pages list, of the second local partition **528** of the persistent memory (B) **526** of the node (B) **524**. Data is also actively mirrored from the second local partition **528** of the persistent memory (B) **526** of the node (B) **524** to the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502** (e.g., data written to a particular page within the second local partition **528** will be mirrored to a corresponding same page within the first remote partition **508**).

(107) Once control of the aggregate (B1) **542** has successfully been transferred from the node (A) **502** to the node (B) **524**, the node (A) **502** may initiate transfer of control of the aggregate (B2) **540** to the node (B) **524**, as illustrated by FIGS. 5F and 5G. Transfer of control of the aggregate (B2) **540** may be performed in a similar manner as how control of the aggregate (B1) **542** was transferred to the node (B) **524**. In particular, the node (A) **502** may allocate a second portion of available free storage space within the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502** as second allocated pages that are allocated/reserved for the aggregate (B2) **540**. The second allocated pages comprise free pages within the first remote partition **508** that are reserved for subsequent use to store data of the aggregate (B2) **540** by the node (A) **502**. The second allocated pages will be different than the allocated pages that were allocated/reserved for the aggregate (B1) **542**, and thus data corruption is avoided because the node (A) **502** will not allocate the same free page for use by the aggregate (B2) **540** (e.g., if control of the aggregate (B2) **540** is retained by the node (A) **502** for serving subsequent I/O directed to the aggregate (B2) **540**) as a corresponding free page having a same page block number that is allocated by the node (B) **524** to store data of the aggregate (B1) **542**.

(108) The node (A) **502** updates metadata information associated with the aggregate (B2) **540** as updated metadata information (B2) **560** for the aggregate (B2) **540**. The updated metadata information (B2) **560** comprises the second allocated pages (e.g., page block numbers of the second allocated pages) tagged with a second identifier of the aggregate (B2) **540**. In this way, the second allocated pages are reserved for subsequent allocation to store data of the aggregate (B2) **540** within the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502**. The allocated pages are removed from the free pages list maintained by the node (A) **502** for the first remote partition **508**. The updated metadata information (B2) **560** is mirrored from the node (A) **502** to the node (B) **524**, as illustrated by FIG. 5F. The node (B) **524** constructs a free pages list for the aggregate (B2) **540** based upon the updated metadata information (B2) **560** for the aggregate (B2) **540** by including page block numbers of the second allocated pages tagged with the second identifier within the updated metadata information (B2) **560**. Once the node (B) **524** has reserved the second allocated pages for the aggregate (B2) **540**, the node (A) **502** transitions control of the aggregate (B2) **540** to the node (B) **524** by performing a giveback **570** of the aggregate (B2) **540**, as illustrated by FIG. 5G. In this way, the node (B) **524** actively processes I/O directed to the

aggregate (B2) **540** using the second local partition **528** of the persistent memory (B) **526** of the node (B) **524**, while mirroring data of the I/O to the first remote partition **508** of the persistent memory (A) **504** of the node (A) **502**.

(109) Once the transition of the aggregates has completed, the soft partition of the local partitions and the remote partitions may be removed. The soft partitions (e.g., the allocation/reservation of certain free pages for use by the aggregate (B1) **542** and the allocation/reservation of different free pages for use by the aggregate (B2) **540**) were created in the event control of one of the aggregates did not get transitioned to the node (B) **524**, and thus the node (B) **524** would be serving I/O for that aggregate while the node (A) **502** would be serving I/O for the other aggregate. In that event, each node will now allocate different free pages to store data of the aggregate that node controls, thus avoiding instances of data corruption and loss where each node could otherwise allocate the same free page to store different data of the different aggregates.

(110) In an embodiment, before the transitioning of the aggregates from the node (A) **502** to the node (B) **524** has successfully completed, available free spaces within the persistent memories of the node (A) **502** and the node (B) **524** may be monitored in the event there is a lack of free space for processing an incoming operation. If the incoming operation targets an aggregate (e.g., aggregate (B2) **540** while still controlled by the node (A) **520**) for which there is a lack of available free space (e.g., a lack of free pages allocated to the aggregate (B2) **540** from the first remote partition **508**) to process the incoming operation, then the incoming operation is suspended, such as queued for subsequent processing once adequate resources become available. A pace at which a scavenger process reclaims unused space (e.g., frees pages as available free pages because those pages are no longer referenced by the active file system and/or snapshots of the active file system) may be increased in response to suspending the incoming operation. In this way, available free pages allocated/reserved for the aggregate (B2) **540** may become available more quickly.

(111) Block allocation for persistent memory during aggregate transition has been described, with respect to FIGS. 5A-5G, in relation to the node (A) **502** performing a giveback procedure after the node (B) **524** recovered from the failure **550**. It may be appreciated that block allocation for persistent memory during aggregate transition can be performed in a similar/same manner for other scenarios. For example, block allocation for persistent memory during aggregate transition may be performed during a takeover, such as a planned takeover where a first node is to takeover aggregates one by one from a second node that is still operational. For each aggregate whose control is being transitioned from the second node to the first node, the second node may allocate pages within persistent memory of the second node as being reserved for an aggregate being transitioned to the first node. Metadata information for the allocated pages is updated with an identifier of the aggregate to create updated metadata information that is mirrored to the first node. Once the first node has allocated/reserved the allocated pages from a persistent memory of the first node, the second node transitions control of the aggregate to the first node. In this way, block allocation for persistent memory during aggregate transition may be performed for the takeover procedure in a same/similar manner as for the giveback procedure.

(112) Still another embodiment involves a computer-readable medium **600** comprising processor-executable instructions configured to implement one or more of the techniques presented herein. An example embodiment of a computer-readable medium or a computer-readable device that is devised in these ways is illustrated in FIG. 6, wherein the implementation comprises a computer-readable medium **608**, such as a compact disc-recordable (CD-R), a digital versatile disc-recordable (DVD-R), flash drive, a platter of a hard disk drive, etc., on which is encoded computer-readable data **606**. This computer-readable data **606**, such as binary data comprising at least one of a zero or a one, in turn comprises processor-executable computer instructions **604** configured to operate according to one or more of the principles set forth herein. In some embodiments, the processor-executable computer instructions **604** are configured to perform a method **602**, such as at least some of the exemplary method **400** of FIG. 4, for example. In some embodiments, the processor-

executable computer instructions 604 are configured to implement a system, such as at least some of the exemplary system 500 of FIGS. 5A-5G, for example. Many such computer-readable media are contemplated to operate in accordance with the techniques presented herein.

(113) In an embodiment, the described methods and/or their equivalents may be implemented with computer executable instructions. Thus, in an embodiment, a non-transitory computer readable/storage medium is configured with stored computer executable instructions of an algorithm/executable application that when executed by a machine(s) cause the machine(s) (and/or associated components) to perform the method. Example machines include but are not limited to a processor, a computer, a server operating in a cloud computing system, a server configured in a Software as a Service (SaaS) architecture, a smart phone, and so on. In an embodiment, a computing device is implemented with one or more executable algorithms that are configured to perform any of the disclosed methods.

(114) It will be appreciated that processes, architectures and/or procedures described herein can be implemented in hardware, firmware and/or software. It will also be appreciated that the provisions set forth herein may apply to any type of special-purpose computer (e.g., file host, storage server and/or storage serving appliance) and/or general-purpose computer, including a standalone computer or portion thereof, embodied as or including a storage system. Moreover, the teachings herein can be configured to a variety of storage system architectures including, but not limited to, a network-attached storage environment and/or a storage area network and disk assembly directly attached to a client or host computer. Storage system should therefore be taken broadly to include such arrangements in addition to any subsystems configured to perform a storage function and associated with other equipment or systems.

(115) In some embodiments, methods described and/or illustrated in this disclosure may be realized in whole or in part on computer-readable media. Computer readable media can include processor-executable instructions configured to implement one or more of the methods presented herein, and may include any mechanism for storing this data that can be thereafter read by a computer system. Examples of computer readable media include (hard) drives (e.g., accessible via network attached storage (NAS)), Storage Area Networks (SAN), volatile and non-volatile memory, such as read-only memory (ROM), random-access memory (RAM), electrically erasable programmable read-only memory (EEPROM) and/or flash memory, compact disk read only memory (CD-ROM)s, CD-Rs, compact disk re-writeable (CD-RW)s, DVDs, cassettes, magnetic tape, magnetic disk storage, optical or non-optical data storage devices and/or any other medium which can be used to store data.

(116) Although the subject matter has been described in language specific to structural features or methodological acts, it is to be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described above. Rather, the specific features and acts described above are disclosed as example forms of implementing at least some of the claims.

(117) Various operations of embodiments are provided herein. The order in which some or all of the operations are described should not be construed to imply that these operations are necessarily order dependent. Alternative ordering will be appreciated given the benefit of this description. Further, it will be understood that not all operations are necessarily present in each embodiment provided herein. Also, it will be understood that not all operations are necessary in some embodiments.

(118) Furthermore, the claimed subject matter is implemented as a method, apparatus, or article of manufacture using standard application or engineering techniques to produce software, firmware, hardware, or any combination thereof to control a computer to implement the disclosed subject matter. The term “article of manufacture” as used herein is intended to encompass a computer application accessible from any computer-readable device, carrier, or media. Of course, many modifications may be made to this configuration without departing from the scope or spirit of the

claimed subject matter.

(119) As used in this application, the terms “component”, “module,” “system”, “interface”, and the like are generally intended to refer to a computer-related entity, either hardware, a combination of hardware and software, software, or software in execution. For example, a component includes a process running on a processor, a processor, an object, an executable, a thread of execution, an application, or a computer. By way of illustration, both an application running on a controller and the controller can be a component. One or more components residing within a process or thread of execution and a component may be localized on one computer or distributed between two or more computers.

(120) Moreover, “exemplary” is used herein to mean serving as an example, instance, illustration, etc., and not necessarily as advantageous. As used in this application, “or” is intended to mean an inclusive “or” rather than an exclusive “or”. In addition, “a” and “an” as used in this application are generally be construed to mean “one or more” unless specified otherwise or clear from context to be directed to a singular form. Also, at least one of A and B and/or the like generally means A or B and/or both A and B. Furthermore, to the extent that “includes”, “having”, “has”, “with”, or variants thereof are used, such terms are intended to be inclusive in a manner similar to the term “comprising”.

(121) Many modifications may be made to the instant disclosure without departing from the scope or spirit of the claimed subject matter. Unless specified otherwise, “first,” “second,” or the like are not intended to imply a temporal aspect, a spatial aspect, an ordering, etc. Rather, such terms are merely used as identifiers, names, etc. for features, elements, items, etc. For example, a first set of information and a second set of information generally correspond to set of information A and set of information B or two different or two identical sets of information or the same set of information.

(122) Also, although the disclosure has been shown and described with respect to one or more implementations, equivalent alterations and modifications will occur to others skilled in the art based upon a reading and understanding of this specification and the annexed drawings. The disclosure includes all such modifications and alterations and is limited only by the scope of the following claims. In particular regard to the various functions performed by the above described components (e.g., elements, resources, etc.), the terms used to describe such components are intended to correspond, unless otherwise indicated, to any component which performs the specified function of the described component (e.g., that is functionally equivalent), even though not structurally equivalent to the disclosed structure. In addition, while a particular feature of the disclosure may have been disclosed with respect to only one of several implementations, such feature may be combined with one or more other features of the other implementations as may be desired and advantageous for any given or particular application.

Claims

1. A method comprising: maintaining pages within a persistent memory of a node, wherein the persistent memory includes a local partition utilized to service client I/O; reclaiming unused pages of the persistent memory that are no longer referenced by at least one of an active file system of the node or a snapshot of the active file system, wherein the unused pages are reclaimed as free pages; modifying metadata information with an aggregate identifier of an aggregate to reserve a set of pages as allocated pages available for subsequently storing data of the aggregate; and utilizing the set of pages to store data of the client I/O directed to the aggregate.
2. The method of claim 1, comprising: constructing a list of the set of pages within the local partition that are available to allocate for storing data of the aggregate; and evaluating the list of the set of pages to identify the set of pages to reserve.
3. The method of claim 1, wherein the reserving further comprises: tagging page block numbers of the set of pages to reserve the set of pages.

4. The method of claim 3, comprising: limiting new allocations, of pages, to the set of pages having the page block numbers tagged.
5. The method of claim 1, wherein the reserving further comprises: assigning the aggregate identifier of the aggregate to page block numbers of the set of pages to reserve the set of pages as the allocated pages for use by the aggregate.
6. The method of claim 1, comprising: maintaining a second node as a partner of the node according to a node pairing; and mirroring data between persistent memories of the node and the second node based upon the node pairing.
7. The method of claim 6, comprising: mirroring the metadata information to the second node to cause the second node to build a second set of pages mirroring the set of pages.
8. The method of claim 7, wherein the set of pages have the same page block numbers of pages maintained by the second node.
9. The method of claim 8, comprising: mirroring data between the set of pages of the node and the pages maintained by the second node.
10. A non-transitory machine readable medium comprising instructions for performing a method, which when executed by a machine, causes the machine to: maintain pages within a persistent memory of a node, wherein the persistent memory includes a local partition utilized to service client I/O; reclaim unused pages of the persistent memory that are no longer referenced by at least one of an active file system of the node or a snapshot of the active file system, wherein the unused pages are reclaimed as pages; modify metadata information with an aggregate identifier of an aggregate to reserve a set of pages as allocated pages available for subsequently storing data of the aggregate; and utilize the set of pages to store data of the client I/O directed to the aggregate.
11. The non-transitory machine readable medium of claim 10, wherein the instructions cause the machine to: construct a list of the set of pages within the local partition that are available to allocate for storing data of the aggregate; and evaluate the list of the set of pages to identify the set of pages to reserve.
12. The non-transitory machine readable medium of claim 10, wherein the instructions cause the machine to: tag page block numbers of the set of pages to reserve the set of pages.
13. The non-transitory machine readable medium of claim 12, wherein the instructions cause the machine to: limit new allocations, of pages, to the set of pages having the page block numbers tagged.
14. The non-transitory machine readable medium of claim 10, wherein the instructions cause the machine to: assign the aggregate identifier of the aggregate to page block numbers of the set of pages to reserve the set of pages as the allocated pages for use by the aggregate.
15. A node comprising: a persistent memory; a memory comprising machine executable code for performing a method; and a processor coupled to the memory, the processor configured to execute the machine executable code to cause the processor to: maintain pages within the persistent memory, wherein the persistent memory includes a local partition utilized to service client I/O; reclaim unused pages of the persistent memory that are no longer referenced by at least one of an active file system of the node or a snapshot of the active file system, wherein the unused pages are reclaimed as pages; modify metadata information to reserve a set of pages available for subsequently storing data of an aggregate; utilize the set of pages to store data of the client I/O directed to the aggregate; and mirror the metadata information to a second node to cause the second node to build a second set of pages mirroring the set of pages.
16. The node of claim 15, wherein the machine executable code causes the processor to: construct a list of the set of pages within the local partition that are available to allocate for storing data of the aggregate; and evaluate the list of the set of pages to identify the set of pages to reserve.
17. The node of claim 15, wherein the machine executable code causes the processor to: tag page block numbers of the set of pages to reserve the set of pages.
18. The node of claim 17, wherein the machine executable code causes the processor to: limit new

allocations, of pages, to the set of pages having the page block numbers tagged.

19. The node of claim 16, wherein the machine executable code causes the processor to: assign an aggregate identifier of the aggregate to page block numbers of the set of pages to reserve the set of pages as allocated pages for use by the aggregate.

20. The node of claim 15, wherein the set of pages have the same page block numbers of pages maintained by the second node.
