

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417099
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Constable; Scott D.

Circuitry and methods for informing indirect prefetches using capabilities

Abstract

Systems, methods, and apparatuses for implementing capability-based indirect prefetching are described. In certain examples, a hardware processor comprises an execution circuit to execute an instruction that generates a memory access request for an element in memory via a first capability; a capability management circuit to check the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; a cache; and a prefetch circuit to: prefetch an additional element of the first object from the memory based on the first capability, determine if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access, and prefetch the second element from the memory to the cache based on the additional element being the second capability.

Inventors:	Constable; Scott D. (Portland, OR)
Applicant:	Intel Corporation (Santa Clara, CA)
Family ID:	1000008820458
Assignee:	Intel Corporation (Santa Clara, CA)
Appl. No.:	17/712073
Filed:	April 02, 2022

Prior Publication Data

Document Identifier	Publication Date
US 20230315465 A1	Oct. 05, 2023

Publication Classification

Int. Cl.: G06F9/30 (20180101)

U.S. Cl.:

CPC G06F9/30047 (20130101); G06F2212/1052 (20130101); G06F2212/6028 (20130101)

Field of Classification Search

CPC: G06F (9/3004); G06F (9/30043); G06F (9/30047); G06F (9/30145); G06F (9/35); G06F (9/383); G06F (9/468); G06F (12/0862); G06F (12/0875); G06F (12/1458); G06F (12/1483); G06F (12/1491); G06F (2212/1052); G06F (2212/452); G06F (2212/602); G06F (2212/6026); G06F (2212/6028); G06F (21/6209); G06F (21/6218); G06F (2221/2141)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
3916385	12/1974	Parmar et al.	N/A	N/A
4809160	12/1988	Mahon et al.	N/A	N/A
4821169	12/1988	Sites et al.	N/A	N/A
5809564	12/1997	Craze et al.	N/A	N/A
5845331	12/1997	Carter et al.	N/A	N/A
5860150	12/1998	Chiarot et al.	N/A	N/A
6009503	12/1998	Liedtke	N/A	N/A
6048940	12/1999	Blaedel et al.	N/A	N/A
6247107	12/2000	Christie	711/216	G06F 12/0862
6671699	12/2002	Black et al.	N/A	N/A
6694450	12/2003	Kidder et al.	N/A	N/A
6823433	12/2003	Barnes et al.	N/A	N/A
7401234	12/2007	Case et al.	N/A	N/A
7948496	12/2010	Hansen et al.	N/A	N/A
8356143	12/2012	Bulusu et al.	N/A	N/A
8554984	12/2012	Yano et al.	N/A	N/A
8595442	12/2012	James-Roxby et al.	N/A	N/A
8850573	12/2013	Chen et al.	N/A	N/A
9026866	12/2014	Balasubramanian	N/A	N/A
9390031	12/2015	Durham et al.	N/A	N/A
9436847	12/2015	Durham et al.	N/A	N/A
9652375	12/2016	Stark et al.	N/A	N/A
9734084	12/2016	Bybell et al.	N/A	N/A
10162694	12/2017	Stark et al.	N/A	N/A
11347508	12/2021	Barnes	N/A	N/A
2004/0031030	12/2003	Kidder et al.	N/A	N/A
2004/0158775	12/2003	Shibuya et al.	N/A	N/A
2005/0102294	12/2004	Coldewey	N/A	G06F 8/4442

2005/0193217	12/2004	Case et al.	N/A	N/A
2006/0187941	12/2005	Andersen	N/A	N/A
2006/0256877	12/2005	Szczepanek et al.	N/A	N/A
2006/0256878	12/2005	Szczepanek et al.	N/A	N/A
2006/0265552	12/2005	Davis et al.	N/A	N/A
2007/0055837	12/2006	Rajagopal et al.	N/A	N/A
2008/0209282	12/2007	Lee et al.	N/A	N/A
2009/0271536	12/2008	Tiennot	N/A	N/A
2009/0292977	12/2008	Bradley et al.	N/A	N/A
2010/0162038	12/2009	Hulbert et al.	N/A	N/A
2013/0254485	12/2012	Kannan et al.	N/A	N/A
2013/0318322	12/2012	Shetty et al.	N/A	N/A
2013/0326288	12/2012	Datta et al.	N/A	N/A
2013/0339649	12/2012	Hsu et al.	N/A	N/A
2014/0115283	12/2013	Radovic et al.	N/A	N/A
2014/0281354	12/2013	Tkacik et al.	N/A	N/A
2014/0372698	12/2013	Lee et al.	N/A	N/A
2015/0134933	12/2014	Holm et al.	N/A	N/A
2015/0199276	12/2014	Radhakrishnan et al.	N/A	N/A
2016/0048378	12/2015	Varma	N/A	N/A
2016/0124802	12/2015	Gabor et al.	N/A	N/A
2016/0259682	12/2015	Stark et al.	N/A	N/A
2016/0283300	12/2015	Stark et al.	N/A	N/A
2016/0371139	12/2015	Stark et al.	N/A	N/A
2017/0177429	12/2016	Stark et al.	N/A	N/A
2018/0060250	12/2017	Hildesheim et al.	N/A	N/A
2018/0074715	12/2017	Farmahini-Farahani et al.	N/A	N/A
2020/0004953	12/2019	Lemay et al.	N/A	N/A
2020/0042464	12/2019	Arulraj	N/A	G06F 12/1483
2020/0133863	12/2019	Pusdesris	N/A	G06F 12/0862
2021/0019148	12/2020	Caulfield	N/A	G06F 9/3455
2021/0081530	12/2020	O'Farrell	N/A	N/A
2021/0200546	12/2020	Lemay et al.	N/A	N/A
2021/0294607	12/2020	Abhishek Raja	N/A	N/A
2023/0085143	12/2022	Smith	N/A	N/A
2023/0195461	12/2022	Lemay et al.	N/A	N/A
2023/0315452	12/2022	Constable	N/A	N/A
2023/0315465	12/2022	Constable	N/A	N/A
2023/0315640	12/2022	Constable	N/A	N/A
2024/0411935	12/2023	Botman	N/A	G06F 9/30189
2025/0156182	12/2024	Botman	N/A	G06F 9/30036

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
------------	------------------	---------	-----

0428079	12/1990	EP	N/A
4198716	12/2022	EP	N/A
03-244054	12/1990	JP	N/A
2007/079011	12/2006	WO	N/A

OTHER PUBLICATIONS

Advisory Action from U.S. Appl. No. 11/323,446, dated Apr. 17, 2012, 3 pages. cited by applicant
 Angelo-Oracle. “SPARC M7 Chip—32 cores”, Oracle.com, Aug. 15, 2014. Web. Accessed Dec. 21, 2015. 8 pages. URL: <http://blogs.oracle.com/rajadurai/entry/sparc_m7_chip_32_cores>. cited by applicant

Arm, “Arm® Architecture Reference Manual Supplement Morello for A-profile Architecture”, Document No. DDI0606, Document Version: A.k, 2019-2022, 1294 pages. cited by applicant

Biin: “CPA Architecture Reference Manual”, 1988, 401 pages. cited by applicant

Burow et al., “CUP: Comprehensive User-Space Protection for C/C++”, Session 9: Software Security, ASIACCS'18, Jun. 4-8, 2018, pp. 381-392. cited by applicant

Carr et al., “DataShield: Configurable Data Confidentiality and Integrity”, Asia CCS '17, Apr. 2-6, 2017, pp. 193-204. cited by applicant

Chen et al., “Shreds: Fine-grained Execution Units with Private Memory”, IEEE Symposium on Security and Privacy, 2016, pp. 56-71. cited by applicant

Cheri, “Capability Hardware Enhanced RISC Instructions (CHERI)”, available online at <<https://www.cl.cam.ac.uk/research/security/ctsrd/cheri/>>- , 2010-2019, 2 pages. cited by applicant

Dartmouth, “ELFbac: Runtime Intent-Level ABI-Granular Memory Protection for Linux”, available online at <<https://www.cs.dartmouth.edu/.about.sergey/io/elfbac/>>, retrieved on May 1, 2020, 1 page. cited by applicant

Duarte, “Memory Translation and Segmentation” Aug. 2008, p. 1-7. cited by applicant

Duck et al., “EffectiveSan: Type and Memory Error Detection using Dynamically Typed C/C++”, PLDI'18, Jun. 18-22, 2018, pp. 181-195. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Dec. 30, 2011, 18 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Jan. 17, 2013, 17 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Jul. 16, 2015, 17 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Jul. 28, 2014, 16 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Jul. 7, 2009, 21 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Oct. 15, 2010, 17 pages. cited by applicant

Final Office Action from U.S. Appl. No. 11/323,446, dated Sep. 20, 2013, 14 pages. cited by applicant

Gil et al., “There's a Hole in the Bottom of the C: On the Effectiveness of Allocation Protection”, IEEE Cybersecurity Development (SecDev), Sep. 30-Oct. 2, 2018, 8 pages. cited by applicant

Gove, D., et al., “Detecting memory access errors,” Nov. 25, 2015, 17 pages. cited by applicant

Gretton-Dann et al., “Arm A—Profile Architecture Developments 2018: Armv8.5-A”, available online at <<https://community.arm.com/developer/ip-products/processors/b/processo-rs-ip-blog/posts/arm-a-profile-architecture-2018-developments-armv85a>>, Sep. 17, 2018, 5 pages. cited by applicant

https://courses.cs.washington.edu/courses/cse351/17wi/lectures/CSE351-L02-memory-I_17wi.pdf

Author: Ceze; Title: CSE351: Memory, Data, & Addressing I, Date: Winter, 2017 (Year: 2017).

cited by applicant

Intel, “Intel® 64 and IA-32 Architectures Software Developer's Manual, vol. 3A: System Programming Guide, Part 1”, Order No. 253668-060US, Sep. 2016, 468 pages. cited by applicant

International Preliminary Report on Patentability for Application No. PCT/US2006/048940, dated Jul. 1, 2008, 8 pages. cited by applicant

International Search Report and Written Opinion for Application No. PCT/US2006/048940, dated Sep. 25, 2007, 15 pages. cited by applicant

International Search Report and Written Opinion for Application No. PCT/US2016/063211, dated Mar. 7, 2017, 11 pages. cited by applicant

Introduction to SPARC M7 and Silicon Secured Memory (SSM), Retrieved from https://swisdev.oracle.com/_files/What-Is-SSM.html on Apr. 7, 2016, 2 pages. cited by applicant

Kwon et al., “Low-Fat Pointers: Compact Encoding and Efficient Gate-Level Implementation of Fat Pointers for Spatial Safety and Capability-based Security”, Proceedings of the 2013 ACM SIGSAC Conference on Computer & Communications Security, 2013, pp. 721-732. cited by applicant

Liljestrand et al., “PAC it up: Towards Pointer Integrity using ARM Pointer Authentication”, Cornell University, Nov. 22, 2018, 21 pages. cited by applicant

LogMeIn Support, “What is Privilege Separation in SSH?”, available online at <https://web.archive.org/web/20191218064501/https://help.logmein.com/articles/en_US/FAQ/What-is-Privilege-Separation-in-SSH-en1>, Dec. 18, 2019, 1 page. cited by applicant

M. Rutland, “ARMv8.3 Pointer Authentication”, Linux Security Summit, Sep. 14, 2017, 24 slides. cited by applicant

Mahon M.J., et al., “Hewlett-Packard Precision Architecture: The Processor,” Hewlett-Packard Journal, Aug. 1986, 19 pages. cited by applicant

McIlroy et al., “Spectre is Here to Stay: An Analysis of Side-Channels and Speculative Execution”, Cornell University, Feb. 15, 2019, pp. 1-26. cited by applicant

Menon et al., “Shakti-T: A RISC-V Processor with Light Weight Security Extensions”, Conference: the Hardware and Architectural Support for Security and Privacy, Jun. 25, 2017, 9 pages. cited by applicant

Miller, Matt, “Trends, Challenges, and Strategic Shifts in the Software Vulnerability Mitigation Landscape”, Microsoft Security Response Center (MSRC), Feb. 7, 2019, 32 pages. cited by applicant

Min R., et al., “Improving Performance of Large Physically Indexed Caches by Decoupling Memory Addresses From Cache Addresses,” IEEE Transactions on Computers, vol. 50, No. 11, Nov. 2001, pp. 1191-1201. cited by applicant

Nagarakatte et al., “CETS: Compiler-Enforced Temporal Safety for C”, ISMM 10, Jun. 5-6, 2010, pp. 31-40. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Apr. 23, 2010, 16 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Jan. 15, 2015, 18 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Jun. 22, 2012, 16 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Jun. 7, 2013, 18 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Mar. 15, 2011, 16 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Mar. 20, 2014, 15 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Nov. 5, 2015, 5 pages. cited by applicant

Non-Final Office Action from U.S. Appl. No. 11/323,446, dated Oct. 7, 2008, 17 pages. cited by applicant

Notice of Allowance from U.S. Appl. No. 11/323,446, dated Mar. 14, 2016, 5 pages. cited by applicant

Osvik et al., “Cache Attacks and Countermeasures: the Case of AES (Extended Version)”, Lecture Notes in Computer Science 2005, Nov. 20, 2005, 25 pages. cited by applicant

Qualcomm, “Pointer Authentication on ARMv8.3: Design and Analysis of the New Software Security Instructions”, Jan. 2017, 12 pages. cited by applicant

Serebryany et al., “AddressSanitizer: A Fast Address Sanity Checker”, 2012 Usenix Annual Technical Conference, Jun. 13-15, 2012, 10 pages. cited by applicant

Serebryany et al., “Memory Tagging and How It Improves C/C++ Memory Safety”, Feb. 2018, 14 pages. cited by applicant

Serebryany et al., “Memory Tagging: How It Improves C/C++ Memory Safety” Google, LLVM Developers' Meeting, Oct. 2018, 29 slides. cited by applicant

Serebryany, Kostya, “Security: ARM Memory Tagging Extension and How It Improves C/C++ Memory Safety”, vol. 44, No. 2, Summer 2019, pp. 12-16. cited by applicant

Suh et al., “Secure Program Execution via Dynamic Information Flow Tracking”, ASPLOS'04, Oct. 9-13, 2004, pp. 85-96. cited by applicant

T. Nyman et al., “HardScope: Thwarting DOP attacks with Hardware-Assisted Run-time Scope Enforcement”, May 2017, pp. 1-20. cited by applicant

The Chromium Projects, “Multi-Process Architecture”, available online at <https://web.archive.org/web/20191030200549/https://www.chromium.org/d-evelopers/design-documents/multi-process-architecture>, Oct. 30, 2019, 2 pages. cited by applicant

Tsampas et al., “Towards Automatic Compartmentalization of C Programs on Capability Machines”, In Proceedings of the International Conference on Foundations of Computer Science, 2017, 14 pages. cited by applicant

Vasilakis et al., “BreakApp: Automated, Flexible Application Compartmentalization”, Network and Distributed Systems Security (NDSS) Symposium, Feb. 18-21, 2018, 15 pages. cited by applicant

Watson et al., “An Introduction to CHERI”, University of Cambridge, Computer Laboratory, Technical Report, No. 941, Sep. 2019, 43 pages. cited by applicant

Watson et al., “Capability Hardware Enhanced RISC Instructions: (CHERI)”, Department of Computer Science and Technology, University of Cambridge, Jan. 2022 (retrieved from <https://www.cl.cam.ac.uk/research/security/ctsrd/cheri/> on Mar. 29, 2022), 3 pages. cited by applicant

Watson et al., “Capability Hardware Enhanced RISC Instructions: CHERI Instruction—Set Architecture (Version 8)”, University of Cambridge, Computer Laboratory, Technical Report, No. 951, Oct. 2020, 590 pages. cited by applicant

Watson et al., “CHERI: A Hybrid Capability-System Architecture for Scalable Software Compartmentalization”, IEEE Symposium on Security and Privacy, 2015, pp. 20-37. cited by applicant

Watson, “Capsicum: Practical capabilities for UNIX”, USENIX Security, 2010, 17 pages. cited by applicant

Wesley et al., “Cornucopia: Temporal Safety for CHERI Heaps”, In Proceedings of the 41st IEEE Symposium on Security and Privacy, 2020, pp. 1507-1524. cited by applicant

Wilkes J., et al., “A Comparison of Protection Lookaside Buffers and the PA-RISC Protection Architecture,” Mar. 1992, Hewlett-Packard, 12 pages. cited by applicant

Xia et al., “CHERIVoke: Characterising Pointer Revocation using CHERI Capabilities for Temporal Memory Safety”, Micro '52, Oct. 2019, pp. 545-557. cited by applicant

Yu et al., “IMP: Indirect Memory Prefetcher”, 48th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), Dec. 5-9, 2015, 13 pages. cited by applicant
Non-Final Office Action, U.S. Appl. No. 17/712,072, May 12, 2025, 18 pages. cited by applicant
Non-Final Office Action, U.S. Appl. No. 17/712,075, Apr. 18, 2025, 9 pages. cited by applicant
Jeon et al., “HexType: Efficient Detection of Type Confusion Errors for C++”, CCS '17, Session K3: Program Analysis, Oct. 2017, pp. 2373-2387. cited by applicant

Primary Examiner: Caldwell; Andrew

Assistant Examiner: Alli; Kasim

Attorney, Agent or Firm: Nicholson De Vos Webster & Elliott LLP

Background/Summary

TECHNICAL FIELD

(1) The disclosure relates generally to electronics, and, more specifically, an example of the disclosure relates to circuitry for informing indirect prefetches using capabilities.

BACKGROUND

(2) A processor, or set of processors, executes instructions from an instruction set, e.g., the instruction set architecture (ISA). The instruction set is the part of the computer architecture related to programming, and generally includes the native data types, instructions, register architecture, addressing modes, memory architecture, interrupt and exception handling, and external input and output (I/O). It should be noted that the term instruction herein may refer to a macro-instruction, e.g., an instruction that is provided to the processor for execution, or to a micro-instruction, e.g., an instruction that results from a processor's decoder decoding macro-instructions.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure is illustrated by way of example and not limitation in the figures of the accompanying drawings, in which like references indicate similar elements and in which:
- (2) FIG. 1 illustrates a block diagram of a hardware processor, including a capability management circuit and a prefetch circuit, coupled to a memory according to examples of the disclosure.
- (3) FIG. 2A illustrates an example format of a capability including a validity tag field, a bounds field, and an address field according to examples of the disclosure.
- (4) FIG. 2B illustrates an example format of a capability including a validity tag field, a permission field, an object type field, a bounds field, and an address field according to examples of the disclosure.
- (5) FIG. 3 illustrates a block diagram of a system comprising a hardware processor, including a plurality of cache levels and a prefetch circuit, coupled to a memory according to examples of the disclosure.
- (6) FIG. 4 illustrates three categories of data structures in memory according to examples of the disclosure.
- (7) FIG. 5 illustrates a memory storing a first capability for a first object “A” and a second capability for a second object “B” according to examples of the disclosure.
- (8) FIG. 6 illustrates an example format for a data structure for capability based prefetching according to examples of the disclosure.

- (9) FIG. 7 is a flow diagram illustrating operations of a method for indirect prefetching based on a capability according to examples of the disclosure.
- (10) FIG. 8 is a flow diagram illustrating operations of another method for indirect prefetching based on a capability according to examples of the disclosure.
- (11) FIG. 9 is a flow diagram illustrating operations of yet another method for indirect prefetching based on a capability according to examples of the disclosure.
- (12) FIG. 10A is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to examples of the disclosure.
- (13) FIG. 10B is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to examples of the disclosure.
- (14) FIG. 11A is a block diagram illustrating fields for the generic vector friendly instruction formats in FIGS. 10A and 10B according to examples of the disclosure.
- (15) FIG. 11B is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 11A that make up a full opcode field according to one example of the disclosure.
- (16) FIG. 11C is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 11A that make up a register index field according to one example of the disclosure.
- (17) FIG. 11D is a block diagram illustrating the fields of the specific vector friendly instruction format in FIG. 11A that make up the augmentation operation field **1050** according to one example of the disclosure.
- (18) FIG. 12 is a block diagram of a register architecture according to one example of the disclosure.
- (19) FIG. 13A is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to examples of the disclosure.
- (20) FIG. 13B is a block diagram illustrating both an exemplary example of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to examples of the disclosure.
- (21) FIG. 14A is a block diagram of a single processor core, along with its connection to the on-die interconnect network and with its local subset of the Level 2 (L2) cache, according to examples of the disclosure.
- (22) FIG. 14B is an expanded view of part of the processor core in FIG. 14A according to examples of the disclosure.
- (23) FIG. 15 is a block diagram of a processor that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to examples of the disclosure.
- (24) FIG. 16 is a block diagram of a system in accordance with one example of the present disclosure.
- (25) FIG. 17 is a block diagram of a more specific exemplary system in accordance with an example of the present disclosure.
- (26) FIG. 18, shown is a block diagram of a second more specific exemplary system in accordance with an example of the present disclosure.
- (27) FIG. 19, shown is a block diagram of a system on a chip (SoC) in accordance with an example of the present disclosure.
- (28) FIG. 20 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to examples of the disclosure.

DETAILED DESCRIPTION

(29) In the following description, numerous specific details are set forth. However, it is understood that examples of the disclosure may be practiced without these specific details. In other instances, well-known circuits, structures, and techniques have not been shown in detail in order not to obscure the understanding of this description.

(30) References in the specification to “one example,” “an example,” “certain examples,” etc., indicate that the example described may include a particular feature, structure, or characteristic, but every example may not necessarily include the particular feature, structure, or characteristic. Moreover, such phrases are not necessarily referring to the same example. Further, when a particular feature, structure, or characteristic is described in connection with an example, it is submitted that it is within the knowledge of one skilled in the art to affect such feature, structure, or characteristic in connection with other examples whether or not explicitly described.

(31) A (e.g., hardware) processor (e.g., having one or more cores) may execute instructions (e.g., a thread of instructions) to operate on data, for example, to perform arithmetic, logic, or other functions. For example, software may request an operation and a hardware processor (e.g., a core or cores thereof) may perform the operation in response to the request. Certain operations include accessing one or more memory locations, e.g., to store and/or read (e.g., load) data. In certain examples, a computer includes a hardware processor requesting access to (e.g., load or store) data and the memory is local (or remote) to the computer. A system may include a plurality of cores, for example, with a proper subset of cores in each socket of a plurality of sockets, e.g., of a system-on-a-chip (SoC). Each core (e.g., each processor or each socket) may access data storage (e.g., a memory). Memory may include volatile memory (e.g., dynamic random-access memory (DRAM)) or (e.g., byte-addressable) persistent (e.g., non-volatile) memory (e.g., non-volatile RAM) (e.g., separate from any system storage, such as, but not limited, separate from a hard disk drive). One example of persistent memory is a dual in-line memory module (DIMM) (e.g., a non-volatile DIMM) (e.g., an Intel® Optane™ memory), for example, accessible according to a Peripheral Component Interconnect Express (PCIe) standard.

(32) Memory may be divided into separate blocks (e.g., one or more cache lines), for example, with each block managed as a single unit for coherence purposes. In certain examples, a (e.g., data) pointer (e.g., an address) is a value that refers to (e.g., points to) the location of data, for example, a pointer may be an (e.g., virtual) address and that data is (or is to be) stored at that address (e.g., at the corresponding physical address). In certain examples, memory is divided into multiple lines, e.g., and each line has its own (e.g., unique) address. For example, a line of memory may include storage for 512 bits, 256 bits, 128 bits, 64 bits, 32 bits, 16 bits, or 8 bits of data, or any other number of bits.

(33) In certain examples, memory corruption (e.g., by an attacker) is caused by an out-of-bound access (e.g., memory access using the base address of a block of memory and an offset that exceeds the allocated size of the block) or by a dangling pointer (e.g., a pointer which referenced a block of memory (e.g., buffer) that has been de-allocated).

(34) Certain examples herein utilize memory corruption detection (MCD) hardware and/or methods, for example, to prevent an out-of-bound access or an access with a dangling pointer. In certain examples, memory accesses are via a capability, e.g., instead of a pointer. In certain examples, the capability is a communicable (e.g., unforgeable) token of authority, e.g., through which programs access all memory and services within an address space. In certain examples, capabilities are a fundamental hardware type that are held in registers (e.g., where they can be inspected, manipulated, and dereferenced using capability instructions) or in memory (e.g., where their integrity is protected). In certain examples, the capability is a value that references an object along with an associated set of one or more access rights. In certain examples, a (e.g., user level) program on a capability-based operating system (OS) is to use a capability (e.g., provided to the program by the OS) to access a capability protected object.

(35) In certain examples of a capability-based addressing scheme, (e.g., code and/or data) pointers are replaced by protected objects (e.g., “capabilities”) that are created only through the use of privileged instructions, for example, which are executed only by either the kernel of the OS or some other privileged process authorized to do so, e.g., effectively allowing the kernel (e.g., supervisor level) to control which processes may access which objects in memory (e.g., without the

need to use separate address spaces and therefore requiring a context switch for an access). Certain examples implement a capability-based addressing scheme by extending memory (e.g., and register) addressing with an additional bit (e.g., writable only in supervisor mode) that indicates that a particular location is a capability, for example, such that all memory accesses (e.g., loads, stores, and/or instruction fetches) must be authorized by a respective capability or be denied. Example formats of capabilities are discussed below in reference to FIGS. 2A and 2B.

(36) In certain examples, a processor uses a prefetcher (e.g., a hardware prefetch circuit) to prefetch data (e.g., program data and/or program instructions), for example, a prefetcher that predicts access addresses and brings the data for those addresses into a cache or caches, e.g., from memory, such as, but not limited to dynamic random access memory (DRAM). In certain examples, accessing memory remains a primary bottleneck to performance in a processor. Prefetching of data is a key technique to solve this problem. In certain examples, prefetching (e.g., hardware memory prefetching) is a performance optimization in which hardware attempts to predict which locations in memory (e.g., machine elements, cache lines, pages, etc.) are likely to be accessed in the near future. In certain examples, these locations are then fetched into a processor's (e.g., central processing unit's (CPU's)) caches, buffers, etc. so that any subsequent accesses will complete at a lower latency, e.g., the lower latency caused by a total time to access that data in the caches, buffers, etc. being (e.g., much) shorter than a total time to service a memory request for that data in the memory.

(37) In certain examples, the prefetcher (e.g., prefetch circuit) lacks high-level semantic information about data structures and program behavior and is to perform guesswork to determine the best candidates for prefetching. If the prefetcher (e.g., prefetch circuit) guesses incorrectly, then frequently accessed (e.g., "hot") data may be displaced and/or evicted from the cache(s) (e.g., buffers, etc.), which can negatively impact performance. If the prefetcher (e.g., prefetch circuit) is too conservative, then opportunities for performance optimization may be missed. To improve the likelihood of true positives and minimize the likelihood of false positives, a prefetcher (e.g., prefetch circuit) may use heuristics that attempt to recognize access patterns exhibited by common design patterns, algorithms, and data structures.

(38) In certain examples, one data structure that may be frequently encountered is a contiguous array of homogeneous elements, for example, where a design pattern (e.g., the iterator pattern) accesses each element of the array in sequence. In certain examples, the iterator searches for an element, aggregates data, and/or modifies one or more of the array elements. Certain hardware prefetchers include built-in heuristics to detect the iterator pattern, e.g., a sequence of memory accesses at a same stride. For example, if the prefetcher (e.g., prefetch circuit) observes that some location X (e.g., a memory address) is accessed, and then a location at $X+N$ bytes is accessed, and then a location at $X+2N$ bytes is accessed, the prefetcher may guess that locations at $X+3N$, $X+4N$, $X+5N$, etc. may be accessed, and may prefetch some (or all) of those locations (e.g., the data at those locations) from memory. However, like all heuristics, this kind of heuristic has limited accuracy. It may not be able to recognize the access pattern if the array is too small (thus, fewer true positives), or it may aggressively prefetch well beyond the end of the array (thus, more false positives).

(39) In certain examples, if the data being prefetched is a pointer to memory or capability to memory, then the pointer or capability can also be used (i) to access the data referenced (e.g., the address referenced) by the pointer or capability (e.g., where the data is data located elsewhere in memory) and/or (ii) to invoke a virtual method on a polymorphic data type, e.g., an object in certain (e.g., C++) programming languages. In certain examples, a hardware prefetcher (e.g., prefetch circuit) pattern-matches the special case where array iteration is used to dereference pointers within the array, and then predicts that future accesses (e.g., at a determined stride) within the array will similarly dereference pointers. Hence, certain prefetchers may not only prefetch memory locations ahead of the iterator, but may also interpret data within those memory locations

as pointers and further prefetch data at the corresponding addresses. In certain examples, this behavior is referred to as an “indirect prefetch” (e.g., even if the “pointer” is an improved pointer in the format of a capability).

(40) In certain examples (e.g., certain processors), the stored bits (e.g., data) may be actual data (e.g., data to be operated on) or a pointer (e.g., an address to already-stored data), but the hardware itself cannot discern between the data and a pointer. Certain examples herein (e.g., certain prefetchers) utilize a hardware (e.g., capability-aware processor) recognized capability (e.g., according to the capability format(s) disclosed herein) to indicate that a value is a capability that points to an address, e.g., such that data stored at the address is prefetched.

(41) Certain instruction set architecture (ISA) designs cause more semantic information to become exposed to hardware. One example is where bounds-checked pointers (e.g., capabilities) are a primitive datatype recognized by the processor (e.g., CPU). Thus, in certain examples, the hardware knows when an element (for example, a single element being a certain (e.g., same) size, e.g., 64 bits) in memory corresponds to a capability, and moreover the hardware also knows the memory bounds (e.g., as a field from the capability) within which the capability is allowed to access. Examples herein utilize one or more of the properties of the capability data type (or any similar technology of other architectures, such as, but not limited to, Intel® Memory Protection Extensions (MPX)) to improve the accuracy of hardware prefetching.

(42) Examples herein improve the accuracy of hardware prefetchers (e.g., prefetch circuits) (e.g., hardware prefetchers that use array-stride detection heuristics) in processors that use a capability-based addressing scheme by also utilizing a capability (e.g., the bounds thereof) to control prefetching of data referenced by the capability, for example, without (or with) adding any additional (e.g., prefetch) field to a format of the capability. Examples herein improve the accuracy of hardware prefetchers (e.g., prefetch circuits) (e.g., hardware prefetchers that use array-stride detection heuristics) in processors that use a hardware implemented capability (e.g., where the capability is provided to a program and/or function by the OS) for obtaining access to a capability protected object via a capability management circuit of the processor by extending the use of the capability (e.g., the bounds thereof) to also control prefetching of data referenced by the capability by a prefetch circuit (e.g., the prefetch circuit separate from the capability management circuit) for example, without (or with) adding any additional (e.g., prefetch) field to a format of the capability. Certain examples herein simplify the hardware that is used to implement a prefetcher (e.g., prefetch circuit) by using capabilities to control prefetching of data referenced by the capability. Examples herein thus provide increased performance and security for a prefetcher (e.g., prefetch circuit), for example, by the microarchitectural innovations disclosed herein that apply a novel performance optimization to an architectural security primitive (e.g., a “capability”).

(43) Examples herein are improvements to the functioning of a processor (e.g., of a computer) itself because they allow a prefetch circuit to utilize a capability (e.g., the bounds thereof) to control prefetching of data referenced by the capability, e.g., where the capability was also prefetched. Examples herein are improvements to the functioning of a processor (e.g., of a computer) itself because they extend the use of a capability (e.g., the hardware enforced bounds thereof) to also control prefetching of data referenced (e.g., bounded) by the (e.g., prefetched) capability by a prefetch circuit. Examples herein allow a prefetch circuit to securely prefetch data under the assumption that a particular (e.g., prefetched) value in memory corresponds to an address pointing elsewhere in memory. Examples herein allow a prefetch circuit to use tagged capabilities (e.g., bounded pointers, “fat” pointers, etc.) to reduce the frequency of false positives, e.g., when performing indirect prefetching on an array, and thus improve performance (e.g., improves the number of instructions per cycle (or clock) (IPC) achieved).

(44) In certain examples (e.g., to reduce the frequency of indirect prefetch false positives during array iteration), a prefetch circuit (and/or prefetch methods) includes (1) augmenting “indirect prefetch” hardware to use memory bounds information encoded in capabilities (e.g., discussed in

reference to FIG. 8) and/or (2) supporting capability-informed “indirect prefetching” during array iteration (e.g., discussed in reference to FIG. 7).

(45) Turning now to the Figures, FIG. 1 illustrates a block diagram of a hardware processor **100** (e.g., core), including a capability management circuit **108** and a prefetch circuit **120**, coupled to a memory **124** according to examples of the disclosure. Although the capability management circuit **108** is depicted within the execution circuit **106**, it should be understood that the capability management circuit can be located elsewhere, for example, in another component of hardware processor **100** (e.g., within fetch circuit **102**) or separate from the depicted components of hardware processor **100**. Although the prefetch circuit **120** is depicted within the memory management circuit **118**, it should be understood that the prefetch circuit **120** can be located elsewhere, for example, in another component of hardware processor **100** (for example, within cache(s) **122**, e.g., a cache controller thereof) or separate from the depicted components of hardware processor **100**.

(46) Depicted hardware processor **100** includes a hardware fetch circuit **102** to fetch an instruction (e.g., from memory **124**), e.g., an instruction that is to request access to a block (or blocks) of memory through a (e.g., non-capability) pointer and/or an instruction that is to request access to a block (or blocks) of memory **124** through a capability **110** to the block (or blocks) of the memory **124**. Depicted hardware processor **100** includes a hardware decoder circuit **104** to decode an instruction, e.g., an instruction that is to request access to a block (or blocks) of memory through a (e.g., non-capability) pointer and/or an instruction that is to request access to a block (or blocks) of memory **124** through a capability **110** to the block (or blocks) of the memory **124**. In certain examples, an instruction is according to a format discussed herein, e.g., in FIGS. 10A-11D. Depicted hardware execution circuit **106** is to execute the decoded instruction, e.g., an instruction that is to request access to a block (or blocks) of memory through a (e.g., non-capability) pointer and/or an instruction that is to request access to a block (or blocks) of memory **124** through a capability **110** to the block (or blocks) of the memory **124**. In certain examples, the execution circuit **106** is to access data in one or more caches (for example, one or more levels of a multiple level cache, e.g., as discussed further in reference to FIG. 3) and/or in the memory **124** (e.g., memory separate from any cache and/or the processor die).

(47) In certain examples, capability management circuit **108** is to, in response to receiving an instruction that is requested for fetch, decode, and/or execution, check if the instruction is a capability instruction or a non-capability instruction (e.g., a capability-unaware instruction), for example, and (i) if a capability instruction, is to allow access to memory **124** storing (1) another capability and/or (2) data and/or instructions (e.g., an object) protected by a capability, and/or (ii) if a non-capability instruction, is not to allow access to memory **124** storing (1) another capability and/or (2) data and/or instructions (e.g., an object) protected by a capability. In certain examples, capability management circuit **108** is to check if an instruction is a capability instruction or a non-capability instruction by checking (i) a field (e.g., opcode) of the instruction (e.g., checking a corresponding bit or bits of the field that indicate if that instruction is a capability instruction or a non-capability instruction) and/or (ii) if a particular register is a “capability” type of register (e.g., instead of a general-purpose data register) (e.g., implying that certain register(s) are not to be used to store a capability or capabilities). In certain examples, capability management circuit **108** is to manage the capabilities, e.g., is to set and/or clear validity tags. In certain examples, capability management circuit **108** is to clear the validity tag of a capability in a register in response to that register being written to by a non-capability instruction.

(48) In certain examples, the source storage location in memory **124** (e.g., virtual address) for a capability **110** is an operand of an (e.g., supervisor level) instruction (e.g., having a mnemonic of LoadCap) that is to load the capability from the memory **124** into register(s) **112** (e.g., into one or more of capability registers **116**). In certain examples, the destination storage location in memory **124** (e.g., virtual address) for a capability **110** is an operand of an (e.g., supervisor level) instruction (e.g., having a mnemonic of StoreCap) that is to store the capability from the register(s) **112** into

memory **124**. In certain examples, the instruction is requested for execution by executing OS code **136** (e.g., or some other privileged process authorized to do so).

(49) In certain examples, capability management circuit **108** is to enforce security properties on changes to capability data (e.g., metadata), for example, for the execution of a single instruction, by enforcing: (i) provenance validity that ensures that valid capabilities can only be constructed by instructions that do so explicitly (e.g., not by byte manipulation) from other valid capabilities (e.g., with this property applying to capabilities in registers and in memory), (ii) capability monotonicity that ensures, when any instruction constructs a new capability (e.g., except in sealed capability manipulation and exception raising), it cannot exceed the permissions and bounds of the capability from which it was derived, and/or (iii) reachable capability monotonicity that ensures, in any execution of arbitrary code, until execution is yielded to another domain, the set of reachable capabilities (e.g., those accessible to the current program state via registers, memory, sealing, unsealing, and/or constructing sub-capabilities) cannot increase.

(50) In certain examples, capability management circuit **108** (e.g., at boot time) provides initial capabilities to the firmware, allowing data access and instruction fetch across the full address space. Additionally, all tags are cleared in memory in certain examples. Further capabilities can then be derived (e.g., in accordance with the monotonicity property) as they are passed from firmware to boot loader, from boot loader to hypervisor, from hypervisor to the OS, and from the OS to the application. At each stage in the derivation chain, bounds and permissions may be restricted to further limit access. For example, the OS may assign capabilities for only a limited portion of the address space to the user software, preventing use of other portions of the address space. In certain examples, capabilities carry with them intentionality, e.g., when a process passes a capability as an argument to a system call, the OS kernel can use only that capability to ensure that it does not access other process memory that was not intended by the user process (e.g., even though the kernel may in fact have permission to access the entire address space through other capabilities it holds). In certain examples, this prevents “confused deputy” problems, e.g., in which a more privileged party uses an excess of privilege when acting on behalf of a less privileged party, performing operations that were not intended to be authorized. In certain examples, this prevents the kernel from overflowing the bounds on a user space buffer when a pointer to the buffer is passed as a system-call argument. In certain examples, these architectural properties of a capability management circuit **108** provide the foundation on which a capability-based OS, compiler, and runtime can implement a certain programming language (e.g., C and/or C++) language memory safety and compartmentalization.

(51) In certain examples, the capability includes an address for data (e.g., instructions and/or data to be operated on by instructions) stored in a single line of memory **124**. In certain examples, the capability includes an address for data (e.g., instructions and/or data to be operated on by instructions) stored in multiple lines of memory **124**. For example, a block of memory may be lines 1 and 2 of data of the (e.g., physical) addressable memory **126** of memory **124** having an address **128** to one (e.g., the first) line (e.g., line 1). In certain examples, a capability includes upper and/or lower bounds of an object protected by a capability, for example, with the address in the capability indicating a single element of multiple elements of a single object (e.g., an array) protected by the capability. Certain examples have a memory of a total size X, where X is any positive integer.

(52) A capability may have different formats and/or fields. In certain examples, a capability is more than twice the width of a native (e.g., integer) pointer type of the baseline architecture, for example, 129-bit capabilities (e.g., including the validity tag) on 64-bit platforms, and 65-bit (e.g., including the validity tag) capabilities on 32-bit platforms. In certain examples, each capability includes an (e.g., integer) address of the natural size for the architecture (e.g., 32 or 64 bit) and also additional metadata (e.g., that is compressed in order to fit) in the remaining (e.g., 32 or 64) bits of the capability. In certain examples, each capability includes (or is associated with) a (e.g., 1-bit) validity “tag” whose value is maintained in registers and memory by the architecture (e.g., by

capability management circuit **108**). In certain examples, each element of the capability contributes to the protection model and is enforced by hardware (e.g., capability management circuit **108**).

(53) In certain examples, capabilities (e.g., one or more fields thereof) themselves are also stored in memory **124**, for example, in data structure **130** (e.g., table) for capabilities. In certain examples, a (e.g., validity) tag **132** is stored in data structure **130** for a capability stored in memory. In certain examples, tags **132** (e.g., in data structure **130**) are not accessible by non-capability (e.g., load and/or store) instructions. In certain examples, a (e.g., validity) tag is stored along with the capability stored in memory (e.g., in one contiguous block).

(54) In certain examples, capability management circuit **108** uses capability-based access control for enforcing memory safety, e.g., and low-overhead compartmentalization (e.g., of objects).

(55) Depicted hardware processor **100** includes one or more registers **112**, for example, general purpose (e.g., data) register(s) **114** (e.g., registers RAX **114A**, RBX **114B**, RCX **114C**, RDX **114D**, etc.) and/or (optional) (e.g., dedicated only for capabilities) capabilities register(s) **116** (e.g., registers CAX **116A**, CBX **116B**, CCX **116C**, CDX **116D**, etc.).

(56) In certain examples, an indication (e.g., name) of the destination register for capability **110** in register(s) **112** is an operand of an (e.g., supervisor level) instruction (e.g., having a mnemonic of LoadCap) that is to load the capability from the memory **124** into register(s) **112**. In certain examples, an indication (e.g., name) of the source register for capability **110** in register(s) **112** is an operand of an (e.g., supervisor level) instruction (e.g., having a mnemonic of StoreCap) that is to store the capability from the register(s) **112** into memory **124**.

(57) Hardware processor **100** includes a coupling (e.g., connection) to memory **124**. In certain examples, memory **124** is a memory local to the hardware processor (e.g., system memory). In certain examples, memory **124** is a memory separate from the hardware processor, for example, memory of a server. Note that the figures herein may not depict all data communication connections. One of ordinary skill in the art will appreciate that this is to not obscure certain details in the figures. Note that a double headed arrow in the figures may not require two-way communication, for example, it may indicate one-way communication (e.g., to or from that component or device). Any or all combinations of communications paths may be utilized in certain examples herein.

(58) In certain examples, hardware processor **100** includes a memory management circuit **118**, for example, to control access (e.g., by the execution unit **106**) to the (e.g., addressable memory **126** of) memory **124**.

(59) As discussed above, in certain examples the execution circuit **106** is to access data in the memory **124** (e.g., memory separate from any cache(s) **122**). However, it may be desirable for prefetch circuit **120** to prefetch certain data from the memory **124** into storage (e.g., cache(s) **122**, registers **112** (e.g., including capabilities register(s) **116**), etc.)) within the processor **100**, e.g., to decrease the time (e.g., waiting) for a memory access request to be serviced in memory **124** by the memory management circuit **118** in comparison to the time to access the storage within the processor (e.g., the cache(s) **122**).

(60) In certain examples, processor **100** (e.g., core) uses a prefetcher (e.g., prefetch circuit **120**) to predict an access (e.g., one or more access addresses) and prefetch (e.g., bring) the data for those one or more addresses into a cache or caches **122** (e.g., data stored in memory **124**, such as, but not limited to dynamic random access memory (DRAM)). In certain examples, one or more prefetch thresholds **121** are indicated for prefetch circuit **120**, for example, a threshold number of accesses to an object (e.g., via a capability) to occur before prefetching data for that object from memory into one or more caches.

(61) In certain examples, processor **100** utilizes capability **110** for access control of data (e.g., protected objects). Examples herein utilize the capability **110** as an input into prefetch circuit **120**, e.g., to prefetch data from memory **124** into cache(s) **122** based at least in part on the capability. Further discussion of this is in reference to FIGS. 4-7. In certain examples, the one or more

elements in memory **124** determined as candidate(s) for prefetching (e.g., into cache(s) **122**) by the prefetch circuit **120** based on an input of a capability are granted corresponding access in the memory **124** for the prefetching by the capability management circuit **108**, e.g., such that prefetching of those element(s) is allowed even though the element(s) are protected by the (e.g., same) capability. In certain examples, the prefetch circuit **120** sends a value to the capability management circuit **108** to indicate which capability that the prefetch corresponds to so that the capability management circuit **108** allows access to the element(s) (e.g., data) protected by that same capability, e.g., to keep the prefetch circuit **120** from allowing prefetch access to element(s) (e.g., data) protected by a different capability.

(62) In certain examples, the prefetching techniques discussed herein are implemented (e.g., entirely) in the frontend of a processor (e.g., with the fetch circuit **102** and decoder circuit **104**) (e.g., before the execution circuit).

(63) In certain examples (e.g., where the prefetch circuit **120** is proximal to the capability management circuit **108**), the capability management circuit **108** is to send information about a load using a capability (e.g., where this differs from a load of a capability itself) to prefetch circuit **120** and this information would become an input to the prefetching discussed herein (e.g., the operations shown in FIG. 7).

(64) In certain examples, the decoder circuit **104** decodes an instruction (e.g., an instruction for a load operation using a capability) into additional “load hint” micro-operation(s) that have the (e.g., physical) load address and capability (e.g., capability bounds) as operands. In certain examples, processor **100** uses load hint micro-operations, but instead of only using the physical address, the load hint micro-operations(s) and are extended to additionally encode the capability (e.g., capability bounds). In certain of these examples, the prefetch circuit **120** does not necessarily need to play a role, e.g., it would depend on how a given processor executes these load hint micro-ops.

(65) In certain examples, processor **100** (e.g., core) uses a prefetcher (e.g., prefetch circuit **120**) to, for prefetched data, further prefetch any data referenced by a capability in the prefetched data (e.g., where the further data that is to be prefetched is also stored at a memory address in memory **124**) into the processor **100** (for example, a core, e.g., a cache or caches **122** thereof). In certain examples, a capability (for example, according to one or more of the format(s) discussed herein, e.g., in FIGS. 2A-2B) is prefetched into processor **100** (e.g., prefetch circuit **120** and/or cache(s) **122**) and, in response to the hardware (e.g., capability management circuit **108**) detecting the prefetched item is a capability (e.g., that points to other data), the prefetch circuit is to issue an additional memory request (or requests) for the one or more further elements based on that prefetched capability (e.g., the one or more further elements protected by that prefetched capability). In certain examples, the hardware (e.g., capability management circuit **108**) detects that prefetched data is a capability by looking up if that prefetched data is associated with (e.g., includes) a (e.g., set) validity tag. In certain examples, the look-up of a validity tag depends on a processor's implementation of tags, e.g., where certain processors utilize tag controller(s) and tag cache(s). In certain examples, a processor reserves a contiguous segment within each data cache (e.g., L1-L3 caches in FIG. 3) to store validity tags. In certain examples, the i.sup.th (e.g., 128-bit) data entry in the data cache has a corresponding validity bit at the i.sup.th offset within the aforementioned reserved contiguous segment, such that when this bit is set, then the processor deduces that the i.sup.th data entry is a valid capability.

(66) In certain examples, processor **100** (e.g., prefetch circuit **120**) utilizes data structure **134** for capability based prefetching. An example of the data structure **134** is discussed herein in reference to FIG. 6.

(67) FIG. 2A illustrates an example format of a capability **110** including a validity tag **110A** field, a bounds **110B** field, and an address **110C** (e.g., virtual address) field according to examples of the disclosure.

(68) In certain examples, the format of a capability **110** includes one or any combination of the

following. A validity tag **110A** where the tag tracks the validity of a capability, e.g., if invalid, the capability cannot be used for load, store, instruction fetch, or other operations. In certain examples, it is still possible to extract fields from an invalid capability, including its address. In certain examples, capability-aware instructions maintain the tag (e.g., if desired) as capabilities are loaded and stored, and as capability fields are accessed, manipulated, and used. A bounds **110B** that identifies the lower bound and/or upper bound of the portion of the address space to which the capability authorizes access (e.g., loads, stores, instruction fetches, or other operations). An address **110C** (e.g., virtual address) for the address of the capability protected data (e.g., object).

(69) In certain examples, the validity tag **110A** provides integrity protection, the bounds **110B** limits how the value can be used (e.g., for example, for memory access), and/or the address **110C** is the memory address storing the corresponding data (or instructions) protected by the capability.

(70) FIG. 2B illustrates an example format of a capability **110** including a validity tag **110A** field, a permission(s) **110D** field, an object type **110E** field, a bounds **110B** field, and an address **110C** field according to examples of the disclosure.

(71) In certain examples, the format of a capability **110** includes one or any combination of the following. A validity tag **110A** where the tag tracks the validity of a capability, e.g., if invalid, the capability cannot be used for load, store, instruction fetch, or other operations. In certain examples, it is still possible to extract fields from an invalid capability, including its address. In certain examples, capability-aware instructions maintain the tag (e.g., if desired) as capabilities are loaded and stored, and as capability fields are accessed, manipulated, and used. A bounds **110B** that identifies the lower bound and/or upper bound of the portion of the address space to which the capability authorizes access (e.g., loads, stores, instruction fetches, or other operations). An address **110C** (e.g., virtual address) for the address of the capability protected data (e.g., object).

Permissions **110D** include a value (e.g., mask) that controls how the capability can be used, e.g., by restricting loading and storing of data and/or capabilities or by prohibiting instruction fetch. An object type **110E** that identifies the object, for example (e.g., in a (e.g., C++) programming language that supports a “struct” as a composite data type (or record) declaration that defines a physically grouped list of variables under one name in a block of memory, allowing the different variables to be accessed via a single pointer or by the struct declared name which returns the same address), a first object type may be used for a struct of people's names and a second object type may be used for a struct of their physical mailing addresses (e.g., as used in an employee directory). In certain examples, if the object type **110E** is not equal to a certain value (e.g., -1), the capability is “sealed” (with this object type) and cannot be modified or dereferenced. Sealed capabilities can be used to implement opaque pointer types, e.g., such that controlled non-monotonicity can be used to support fine-grained, in-address-space compartmentalization.

(72) In certain examples, permissions **110D** include one or more of the following: “Load” to allow a load from memory protected by the capability, “Store” to allow a store to memory protected by the capability, “Execute” to allow execution of instructions protected by the capability, “LoadCap” to load a valid capability from memory into a register, “StoreCap” to store a valid capability from a register into memory, “Seal” to seal an unsealed capability, “Unseal” to unseal a sealed capability, “System” to access system registers and instructions, “BranchSealedPair” to use in an unsealing branch, “CompartmentID” to use as a compartment ID, “MutableLoad” to load a (e.g., capability) register with mutable permissions, and/or “User[N]” for software defined permissions (where N is any positive integer greater than zero).

(73) In certain examples, the validity tag **110A** provides integrity protection, the permission(s) **110D** limits the operations that can be performed on the corresponding data (or instructions) protected by the capability, the bounds **110B** limits how the value can be used (e.g., for example, for memory access), the object type **110E** supports higher-level software encapsulation, and/or the address **110C** is the memory address storing the corresponding data (or instructions) protected by the capability.

(74) In certain examples, a capability (e.g., value) includes one or any combination of the following fields: address value (e.g., 64 bits), bounds (e.g., 87 bits), flags (e.g., 8 bits), object type (e.g., 15 bits), permissions (e.g., 16 bits), tag (e.g., 1 bit), global (e.g., 1 bit), and/or executive (e.g., 1 bit). In certain examples, the flags and the lower 56 bits of the “capability bounds” share encoding with the “capability value”.

(75) In certain examples, the format of a capability (for example, as a pointer that has been extended with security metadata, e.g., bounds, permissions, and/or type information) is the same bit width (other than the validity tag) as a native pointer, e.g., a capability having a (e.g., 64-bit) address (e.g., pointer) combined with bounds information and/or other metadata that can be used for enforcing other fine-grained access control policies. In certain examples, a capability includes a (e.g., 1-bit) validity tag that is guarded within processor-protected (e.g., “out of band”) memory, for example, such that this tag serves as a (e.g., tamper-proof) type indicator that indicates that a given value in memory is a capability (e.g., in contrast with that value being data or a non-capability pointer).

(76) FIG. 3 illustrates a block diagram of a system **300** comprising a hardware processor **100**, including a plurality of cache levels and a prefetch circuit **120**, coupled to a memory **124** according to examples of the disclosure.

(77) A processor may include a prefetch circuit, e.g., the processors discussed below. FIG. 3 illustrates an example of multiple processor cores (core A **100-A** and core B **100-B**) and multiple levels of caches (e.g., L1, L2, and L3), e.g., in a cache coherency hierarchy. Although two cores are depicted, a single or more than two cores may be utilized. Although multiple levels of cache are depicted, a single, or any number of caches may be utilized. Cache(s) may be organized in any fashion, for example, as a physically or logically centralized or distributed cache.

(78) In an example, a processor **100**, such as a processor or processors including the processor cores illustrated in the Figures, or any other processor, may include one or more caches. FIG. 3 illustrates an example of a three level (e.g., levels 1 (L1), 2 (L2), and 3 (L3)/last level cache (LLC)) cache hierarchy. In certain examples, multiple cores (core A **100-A** and B **100-B**) are of a single processor **100**. A core (e.g., core A and core B) may include the components of a processor to execute instructions. An un-core may include all logic circuitry not in a core. A processor core (e.g., core A) may include components such as a level 1 instruction cache (L1-INST.) **302** and a level 1 data cache (L1-DATA) **304**. A core may include components such as an address generation unit (AGU) **308**, translation lookaside buffer (TLB) **310**, and a level 2 cache (L2) **306**. A core may or may not share a cache with other cores, for example, core A and core B may share some level(s) of cache and not share other level(s) of cache, e.g., core A and core B may share the level 3 cache (L3) **312** but not their L2 **306** or L1 (**302,304**) caches. A core may include any combination of these components or none of these components. Processor **100** (e.g., core A and core B) may access (e.g., load and store) data in the (e.g., system) memory **124**, e.g., as indicated by the arrow. In one example, the (e.g., system) memory **124** communicates with the core over a bus and/or interconnect, e.g., at a slower access and/or cycle time than the core accessing cache (e.g., cache on the processor **100**).

(79) An address generation unit (e.g., AGU **308**), for example, address computation unit (ACU), may refer to a circuit inside a processor (e.g., a core) that calculates addresses used to access memory (e.g., system memory), for example, to allow the core to access the system memory. In one example, the AGU takes data values (e.g., register value and addresses mentioned in an instruction) as an input and outputs the (e.g., virtual) addresses for that stream. A core may include an execution circuit (e.g., execution unit) to perform arithmetic operations, such as addition, subtraction, modulo operations, or bit shifts, for example, utilizing an adder, multiplier, shifter, rotator, etc. thereof.

(80) A translation lookaside buffer (e.g., TLB **310**) may convert a virtual address to a physical address (e.g., of the system memory). A TLB may include a data table to store (e.g., recently used) virtual-to-physical memory address translations, e.g., such that the translation does not have to be

performed on each virtual address present to obtain the physical memory address. If the virtual address entry is not in the TLB, a processor may perform a page walk to determine the virtual-to-physical memory address translation.

(81) Prefetch circuit **120** (for example, one or more parallel instances of prefetch circuits **120-1** to **120-N**, e.g., where N is any positive integer greater than one) may be a separate functional circuit, for example, not utilizing a functional circuit (e.g., execution circuit (e.g., Arithmetic Logic Unit (ALU)), AGU, TLB, etc.) of a core. Prefetch circuit may be part of a processor (e.g., but separate from a core(s)). Prefetch circuit may communicate with the core(s) of the processor, e.g., via communication resources, such as, but not limited to, a ring network. Processor **100** may communicate with the (e.g., system) memory **124** and/or caches (e.g., L1, L2, or L3 in FIG. 1) via a memory controller (e.g., as part of the processor) and/or an interconnect. Prefetch circuit **120** may output the system memory addresses to a memory controller (e.g., memory management circuit **118** in FIG. 1 or memory unit **1370** in FIG. 13B) of processor **100**. In certain examples, each of core A **100-A** and core B **100-B** includes execution circuitry, e.g., one or more (e.g., a plurality of) execution circuits (e.g., as instance(s) of execution circuit **106** in FIG. 1).

(82) In certain examples, prefetch circuit **120** includes circuitry (e.g., hardware logic circuitry) to perform the prefetching (e.g., indirect prefetching) discussed herein, e.g., based on a capability.

(83) In certain examples, prefetching based on a capability is utilized for one or more of certain categories of (for example, (e.g., logically) contiguous) data structures in memory.

(84) FIG. 4 illustrates three categories (**402**, **404**, and **406**) of data structures in memory according to examples of the disclosure. In certain examples, a first category **402** is a single (e.g., machine) element of a primitive type, for example, (e.g., signed (S) or unsigned (U)) (e.g., 64-bit) integer (INT), pointer, capability, single-precision floating-point value, double-precision floating-point value, etc. In certain examples, a second category **404** is a heterogeneous structure comprised of one or more elements of varying types, e.g., including capabilities and/or pointers to other structures located elsewhere in memory. Examples include a single node in a linked list or in a binary tree. In certain examples, a third category **406** is a contiguous array of homogeneous elements (e.g., an array of values (e.g., integers or pointers)) or an array of (e.g., heterogeneous) structures (e.g., a combination of integers, capabilities, and/or pointers).

(85) FIG. 5 illustrates a memory **124** storing a first capability **504** for a first object “A” **504-1** to **504-M** (e.g., where M is a positive integer greater than one) and a second capability **510** for a second object “B” **510-1** to **510-M** (e.g., where N is a positive integer greater than one) according to examples of the disclosure. In certain examples, an array **500** includes one or more elements stored therein that are addresses, e.g., capabilities for those addresses. In certain examples, array **500** is a third object “C” having eight elements indexed 0 to 7.

(86) In certain examples, the array **500** itself is indicated by pointer(s) and/or capability(ies), for example, a non-capability pointer “I” or capability “I” that identifies one of the elements of the array **500**, e.g., via a capability's address field (e.g., as shown in FIG. 2A or 2B). In the depicted example in FIG. 5, “I” can either be (a) a capability that points to C[4] and protects object “C”, or (b) a raw (e.g., non-capability) pointer that points to C[4], e.g., where 4 is the index into the elements of the object. In certain examples, one or more of the elements of array **500** are prefetched (e.g., as indicated by a capability “I” or a raw pointer “I” for one of those elements), and those one or more elements of array **500** are checked to determine if any are capabilities (e.g., via checking whether a corresponding validity tag is set that indicates that data is a valid capability), and if so, then, causing an additional prefetch of the memory location(s) indicated by that capability, and/or if not, not causing an additional prefetch that assumes those element(s) of the array **500** are memory location(s).

(87) In certain examples, at least second element **504** and fifth element **510** (e.g., because they are capabilities) are prefetched from memory **124** into a processor (e.g., processor **100** in the figures) (e.g., into at least the prefetch circuit **120** and/or cache(s) **122** thereof). In certain examples, the

array is protected by a capability "I" that indicates array "C" includes eight elements, so each of lines 1-8 are prefetched into a processor (e.g., processor **100** in the figures) (e.g., into at least the prefetch circuit **120** and/or cache(s) **122** thereof).

(88) Although the following paragraph references data as stored in memory **124**, it should be understood that the array **500** is prefetched into a processor (e.g., processor **100** in the figures) (e.g., into at least the prefetch circuit **120** and/or cache(s) **122** thereof) in certain examples. In certain examples, a first element **502** (e.g., line) from memory stores a pointer or data that the hardware (e.g., capability management circuit **108**) detects is not a capability, a second element **504** (e.g., line) from memory stores a capability that the hardware (e.g., capability management circuit **108**) detects as a capability (e.g., not data and not a non-capability pointer), a third element **506** (e.g., line) from memory stores a pointer or data that the hardware (e.g., capability management circuit **108**) detects is not a capability, a fourth element **508** (e.g., line) from memory stores a pointer or data that the hardware (e.g., capability management circuit **108**) detects is not a capability, and/or a fifth element **510** (e.g., line) from memory stores a capability that the hardware (e.g., capability management circuit **108**) detects as a capability (e.g., not data and not a non-capability pointer).

(89) In certain examples, a capability includes an address that points to a single one of multiple elements of an object protected by that capability, for example, with the capability "A" from second element **504** (e.g., line) from memory pointing to first element **504-1** (e.g., line) of object A and the capability "B" from fifth element **510** (e.g., line) from memory pointing to a second element **510-2** (e.g., line) of object B.

(90) In certain examples, a prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 2 at **504**, that it is a capability (e.g., first capability "A"), and, based on that being a capability (e.g., instead of a non-capability value), then prefetch one or more elements of object "A" that are protected by the first capability "A", e.g., to prefetch one or more (e.g., all) element(s) **504-1** to **504-M** of object "A" identified by the first capability "A" (for example, via its upper and/or lower bounds field, e.g., bounds **110B** in FIGS. 2A and 2B). In certain examples, the prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 2, that it is a capability (e.g., first capability "A") pointing to an (e.g., single) element in object A and prefetch that element. In certain examples, a prefetcher (e.g., prefetch circuit **120**) detects that there are one or more additional elements before and/or after the "pointed to" element of the capability and causes them to be prefetched, for example, if line 1 **504-1** of object A is pointed to by the capability A, the prefetcher is then to prefetch line 2 **504-2** up to line M **504-M** (e.g., from memory **124** into a cache).

(91) In certain examples, a prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 5 at **510**, that it is a capability (e.g., second capability "B"), and, based on that being a capability (e.g., instead of a non-capability value), then prefetch one or more elements of object "B" that are protected by the second capability "B", e.g., to prefetch one or more (e.g., all) element(s) **510-1** to **510-N** of object "B" identified by the first capability "B" (for example, via its upper and/or lower bounds field, e.g., bounds **110B** in FIGS. 2A and 2B). In certain examples, the prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 5, that it is a capability (e.g., first capability "B") pointing to an (e.g., single) element in object B and prefetch that element. In certain examples, a prefetcher (e.g., prefetch circuit **120**) detects that there are one or more additional elements before and/or after the "pointed to" element of the capability and causes them to be prefetched, for example, if line 1 **510-2** of object B is pointed to by the capability B, the prefetcher is then to prefetch line 1 **510-1** and/or line 3 up to line N **510-N** (e.g., from memory **124** into a cache).

(92) In certain examples, a data structure (e.g., data structure **134**) is utilized by the prefetch circuit to control prefetching. In certain examples, the data structure **134** is within the prefetch circuit **120**. In certain examples, the data structure **134** is not within the prefetch circuit **120**, e.g., depending on a processor's implementation. For example, if the processor has an execution circuit **106** that executes load hints, then the data structure **134** could be located within or adjacent to that execution

circuit **106**. For example, suppose the prefetch circuit **120** consumes a capability load hint with address/offset value “x” and capability metadata “b” (e.g., including bounds, permissions, etc.), thus in certain examples, the execution circuit **106** uses “b” and “x” as inputs to **706** and **708**, respectively, to perform the operations in FIG. 7 and issue loads (e.g., through the prefetch circuit or through some other implementation-dependent circuit(s)).

(93) In other examples, at least second element **504** and fifth element **510** (e.g., because they are capabilities) are prefetched from memory **124** into a processor (e.g., processor **100** in the Figures) (e.g., into at least the prefetch circuit **120** and/or cache(s) **122** thereof). In certain of these examples, the array “C” is accessed either using a capability and the processor does not use the capability bounds to inform the prefetching of elements of “C” (but these elements may still be prefetched by some other heuristic), or “C” is accessed using a raw (e.g., non-capability) pointer and elements of “C” may still be prefetched by some other heuristic. In certain of these examples the data prefetched (which may or may not be within the bounds of “C”) may, e.g., at a stride of “D” bytes, be scanned for valid capabilities. Each scanned location that contains a valid capability may have its protected element(s) be prefetched by the processor in the manner described in the prior paragraph, e.g., but according to the operations in FIG. 8.

(94) FIG. 6 illustrates an example format for a data structure **134** for capability based prefetching according to examples of the disclosure. In certain examples, one or any combination (e.g., all) of the depicted types (shown as columns) may be utilized. In FIG. 6, data structure **134** is shown as a table including a first column for a capability index **602** (e.g., with the value in an entry (shown as a row) indicating which capability (e.g., object protected by that capability) of a plurality of capability, a second column for a most recent offset accessed **604** (for example, a most recent offset (e.g., element by element offset) that was accessed in that capability) (for example, if the most recent access to object “B” in FIG. 5 was line 2 **510-2**, offset **604** is set to a value that identifies line 2 of object “B”), a third column for a most recent offset prefetched **606** (for example, a most recent offset (e.g., element by element offset) that was prefetched for that capability) (for example, if the most recent prefetch for object “B” in FIG. 5 was line 3, offset **606** is set to a value that identifies line 3, e.g., such that additional prefetches are tracked as they occur), a fourth column for a thread identification (ID) value and/or a logical processor ID value **608** (for example, a thread ID value and/or a logical processor ID value that indicate which hardware thread of a plurality of hardware threads and/or logical processor of a plurality of logical processors that created and/or updated the entry, respectively), and/or a fifth column for a valid and/or timeout and/or age value **610** (for example, a valid bit that when set (e.g., to one) indicates the entry (shown as a row) is valid) and when cleared (e.g., to zero) indicates the entry is invalid and thus not to be used and/or a timeout value to indicate a time for the entry so that the entry can be invalidated after the elapse of a time period).

(95) In certain examples, data structure **134** is implemented in hardware, e.g., as part of prefetch circuit and/or capability management circuit. In certain examples, data structure **134** associates each capability in the data structure with the most recent access that has been made using the capability. Such a table can be implemented (e.g., compactly) as a direct lookup table, e.g., as a table where the domain (e.g., a representation of the capability) is used as an index to select the corresponding element of the codomain, the most recent access that was made using the capability.

(96) As permitted by hardware, data structure **134** (e.g., table) may contain $2^{\lceil n \rceil}$ entries, where n is the number of bits (e.g., a proper subset of bits that represent a capability) that are sampled from the capability, for example, where “n” is the number of a subset of bits sampled from the capability or “n” is the hash of a subset of bits sampled from the capability. The example in FIG. 6 has n=3 where, for instance, the lowest-order 3 bits of each capability are used to form the index. In certain examples, a larger value of n will reduce false positive prefetches triggered by collisions within the data structure **134** (e.g., table), but may also consume more die space on a physical chip.

(97) In certain examples, each entry in the data structure **134** (e.g., table) includes m bits (e.g., a proper subset of bits) that are sampled from the most recent offset within the capability bounds that was accessed. For instance, this could be the lowest-order m bits of the offset (in the example in FIG. 6, $m=8$). In certain examples, the total size of the data structure **134** (e.g., as a direct lookup table) is $m \times 2^{\lceil \log_2(n) \rceil}$ bits. In certain examples, a higher value of m will reduce the frequency of false positives due to collisions, but may also cost more die area.

(98) Supporting Capability-Informed “Indirect Prefetching” During Array Iteration

(99) FIG. 7 shows examples operations of the lookup data structure **134** (e.g., table) in conjunction with a memory access through a capability, and that (optionally), for another capability that is prefetched, further prefetches one or more of the memory locations indicated by that prefetched capability. More particularly, FIG. 7 is a flow diagram illustrating operations **700** of a method for prefetching based on a capability according to examples of the disclosure. Some or all of the operations **700** (or other processes described herein, or variations, and/or combinations thereof) are performed under the control of a prefetch circuit and/or capability management circuit as implemented herein and/or one or more computer systems configured with executable instructions and are implemented as code (e.g., executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. The code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising instructions executable by one or more processors. The computer-readable storage medium is non-transitory. In some examples, one or more (or all) of the operations **700** are performed by prefetch circuit **120** and/or capability management circuit **108** of the other figures.

(100) The operations **700** include, at block **702**, starting the operations, e.g., enabling prefetch based on a capability. The operations **700** further include, at block **704**, receiving a memory access request through a capability (e.g., this access being detected at or after execution (e.g., in an execution stage of a pipelined processor) and/or at or after retirement (e.g., in a retirement stage of a pipelined processor)). The operations **700** further include, at block **706**, using the capability (for example, all of the capability bits or a proper subset of the capability bits, e.g., a hash of a plurality of (e.g., all of) capability bits) to index into a data structure for capability based prefetching (e.g., data structure **134**). For example, if the capability corresponds to an index of 001, accessing the second entry in the data structure **134** as depicted in FIG. 6. In certain examples, any sequence of bits that are associated with, or are derived from, the capability is used to index into the data structure for capability based prefetching (e.g., data structure **134**). The operations **700** further include, at block **708**, comparing a current offset from the capability (e.g., current element of a plurality of elements of an object corresponding to the capability) to a saved offset from the data structure (e.g., most recent offset accessed **604** in FIG. 6 for that entry) for capability based prefetching (e.g., data structure **134**) to generate the delta (“D”) of those offsets. The operations **700** further include, at block **710**, checking if the delta (“D”) is zero, and if yes, it is the end of the operations **700** at block **718**, and if no, proceeding to block **712** for prefetching of one or more (e.g., “K”) memory locations of the capability (e.g., not exceeding the bound(s) of the capability) (e.g., at a stride of “D” bytes relative to the current offset). The operations at block **712** may include not prefetching any data (e.g., elements) that were already prefetched (e.g., based on the most recent offset prefetched **606** in FIG. 6 for that entry). The operations **700** further include, at block **714**, for (e.g., each) of the one or more (e.g., “K”) memory locations of the capability that themselves store a capability, prefetching the capability's corresponding memory locations (e.g., element(s) stored in those memory locations) (e.g., and if not a capability, not performing a prefetch that assumes that data from one of the “K” memory locations was a pointer and/or capability) and proceeding to block **716**. The operations **700** further include, at block **716**, updating the (e.g., same entry of the) data structure for capability based prefetching (e.g., data structure **134**) with the current offset, and then proceeding to the end of the operations **700** at block **718**.

(101) Referring to FIG 5 (e.g., where “I” is a non-capability pointer or a capability that identifies one of the elements of the array **500**), as one example, if the prefetch at block **712** is of one or more elements (e.g., all elements) of array **500** in FIG. 5, then, at block **714**, prefetcher (e.g., prefetch circuit **120**) is to: (i) detect, for the prefetch of the line 2 at **504**, that it is a capability (e.g., first capability “A”), and, based on that being a capability (e.g., instead of a non-capability value), then prefetch one or more elements of object “A” that are protected by the first capability “A”, e.g., to prefetch one or more (e.g., all) element(s) **504-1** to **504-M** of object “A” identified by the first capability “A” (for example, via its upper and/or lower bounds field, e.g., bounds **110B** in FIGS. 2A and 2B). In certain examples, the prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 2, that it is a capability (e.g., first capability “A”) pointing to an (e.g., single) element in object A and prefetch that element. In certain examples, a prefetcher (e.g., prefetch circuit **120**) detects that there are one or more additional elements before and/or after the “pointed to” element of the capability and causes them to be prefetched, for example, if line 1 **504-1** of object A is pointed to by the capability A, the prefetcher is then to prefetch line 2 **504-2** up to line M **504-M** (e.g., from memory **124** into a cache); and/or (ii) detect, for a prefetch of the line 5 at **510**, that it is a capability (e.g., second capability “B”), and, based on that being a capability (e.g., instead of a non-capability value), then prefetch one or more elements of object “B” that are protected by the second capability “B”, e.g., to prefetch one or more (e.g., all) element(s) **510-1** to **510-N** of object “B” identified by the first capability “B” (for example, via its upper and/or lower bounds field, e.g., bounds **110B** in FIGS. 2A and 2B). In certain examples, the prefetcher (e.g., prefetch circuit **120**) is to detect, for a prefetch of the line 5, that it is a capability (e.g., first capability “B”) pointing to an (e.g., single) element in object B and prefetch that element. In certain examples, a prefetcher (e.g., prefetch circuit **120**) detects that there are one or more additional elements before and/or after the “pointed to” element of the capability and causes them to be prefetched, for example, if line 1 **510-2** of object B is pointed to by the capability B, the prefetcher is then to prefetch line 1 **510-1** and/or line 3 up to line N **510-N** (e.g., from memory **124** into a cache).

(102) Note that the data structure for capability based prefetching (e.g., data structure **134**) and its behavior can be augmented to further improve the precision. For example, some examples on processors that support simultaneous multithreading (SMT) (e.g., Intel® Hyper-Threading) tag each entry with a bit(s) corresponding to the hardware thread and/or logical processor that created the entry, e.g., such that an (e.g., each) entry is only used and/or updated for actions by that corresponding hardware thread and/or logical processor. In certain examples, entries in the data structure have associated bit(s) corresponding to the lifetime of the entry, for example, a valid bit that is cleared after some duration and/or cleared on a specific event (e.g., a context switch).

(103) In certain examples, the dashed block at **714** in FIG. 7 is equivalent to the loop depicted in FIG. 8 with dashed blocks (blocks **808**, **810**, **812**, **814**, and **816**), but for brevity, it is shown as a single box in FIG. 7). In other examples, the dashed block at **714** and/or dashed blocks in FIG. 8 are implemented with parallel prefetch circuits, e.g., depicted in FIG. 3.

(104) In certain examples, the operations shown in FIG. 7 can be implemented as an extension to hardware prefetchers that are not already enlightened to capabilities, e.g., a prefetcher that does not use a capability as an input (e.g., a “capability unaware” prefetch circuit). In certain examples, the operations shown in FIG. 7 reduce the number of false positives. For example, consider two arrays of pointers A and B that are adjacent in memory. Suppose that a program is iterating through A and dereferencing pointers. In certain examples, prefetcher is unaware of the bounds of A, and therefore may observe capabilities in array B and prefetch their memory references, even though these capabilities will not be used while iterating through A. In contrast, in certain examples, the method shown in FIG. 7 is aware of the bounds of A, and therefore will only perform the indirect prefetching within A, e.g., unless or until the program begins to iterate through B. Certain examples herein thus prevent anything other than a capability from being treated as a memory reference, e.g.,

such that it is not possible for some arbitrary value in memory to be considered as an indirect prefetch (e.g., “pointer prefetch”) candidate and/or used for the indirect prefetch.

(105) Augmenting “Indirect Prefetch” Hardware to Use Memory Bounds Information Encoded in Capabilities

(106) FIG. 8 is a flow diagram illustrating operations **800** of another method for indirect prefetching based on a capability according to examples of the disclosure. Some or all of the operations **800** (or other processes described herein, or variations, and/or combinations thereof) are performed under the control of a prefetch circuit and/or capability management circuit as implemented herein and/or one or more computer systems configured with executable instructions and are implemented as code (e.g., executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. The code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising instructions executable by one or more processors. The computer-readable storage medium is non-transitory. In some examples, one or more (or all) of the operations **800** are performed by prefetch circuit **120** and/or capability management circuit **108** of the other figures.

(107) The operations **800** include, at block **802**, starting the operations, e.g., enabling (e.g., indirect) prefetch based on a capability. The operations **800** include, at block **804**, detecting, via hardware (for example, a processor, e.g., a prefetch circuit of a system), one or more (for example, a plurality greater than a threshold, e.g., where the threshold is selectable in the system) pointer and/or capability dereferences (e.g., at a stride of “D” bytes, where D is a positive integer), for example, detecting that data (e.g., data prefetched in a cache) being accessed (for example, in an array, e.g., at the stride of “D” bytes) is a pointer(s) and/or capability, e.g., in contrast to a value that does not point to a different memory location. The operations **800** further include, at block **806**, prefetching, via hardware (for example, the processor, e.g., the prefetch circuit of a system), one or more (e.g., “K”) memory locations (e.g., at the stride of “D” bytes). The operations **800** further include, at block **808**, setting an index “i” to an initial value (e.g., 1) to indicate that an “element to be checked” is the initial element of the prefetched (e.g., “K”) memory location(s). The operations **800** further include, at block **810**, checking if the indexed element (e.g., starting at an index value of 1) of the prefetched (e.g., “K” memory locations) includes (e.g., is) a capability, and (i) if no at block **810**, incrementing the index value (e.g., by one) at block **814** and, at block **816**, checking if all of the “K” memory locations have been checked (e.g., is “i” greater than “K”), and, if yes at block **816**, proceeding to the end of the operations **800** at block **818**, and if no at block **816**, returning to block **810**; and if (ii) yes at block **810**, prefetching the memory location(s) referenced by the capability (e.g., prefetching the single element that is addressed in the capability or a plurality (e.g., all of) the elements protected by the capability (e.g., within the bounds of the capability).

(108) In certain examples, one or more of the operations **800** are performed by a prefetch circuit (e.g., prefetch circuit **120**) and/or a capability management circuit (e.g., capability management circuit **108**). In one example, the capability check(s) at block **810** is performed by a capability management circuit, for example, and those results are passed to a prefetch circuit (e.g., to cause the prefetch(es) at block **812**.

(109) Referring to FIG. 5 (e.g., where “I” is a capability that identifies one of the elements of the array **500** and protects the array as object “C”), as one example, if the prefetch at block **806** is of one or more elements of array **500** in FIG. 5, then, at block **810**, the checking (e.g., by a capability management circuit) is if any of the elements (e.g., lines 1-8 in array “C”) are capabilities, e.g., and where prefetched element 2 (e.g., index “1”) **504** (e.g., one or more cache lines) is a capability (for object “A”), causing, at block **812**, the prefetch of one or more elements (e.g., one or more cache lines) **504-1** to **504-M** of object “A” and/or where prefetched element 5 (e.g., index “4”) **510** (e.g., one or more cache lines) is a capability (for object “B”), causing, at block **812**, the prefetch of one

or more elements (e.g., one or more cache lines) **510-1** to **510-M** of object “B”.

(110) In certain examples, the check at block **810** does not actually have access to modify the capability (e.g., validity tag thereof), for example, it is only to check that there is an active capability (e.g., indicated via a validity tag set to “valid”) or not.

(111) In certain examples, a processor (e.g., CPU) implements indirect prefetch hardware with detection heuristics that identify pointer dereferences at a stride. In certain examples, the processor (e.g., prefetch circuit) detects that a sequence of (e.g., 64-bit) elements at a stride of a certain number of (e.g., 64) bytes have all been used to dereference memory, and the processor (e.g., prefetch circuits) is then to trigger, for example, a certain number of (e.g., 32) prefetches at the stride (e.g., of 64 bytes) beyond the most recently detected dereference. Thus, certain examples herein provide for a detect-and-prefetch system and method to be extended to take advantage of memory bounds information encoded in a capability (including bounded pointer, “fat” pointer, Intel® MPX, etc.). In certain examples, an (e.g., array) prefetch circuit is modified (e.g., via the dotted structures in FIG. **8**) to conduct indirect prefetching that is informed by a capability (e.g., an array bounds encapsulated in the capability).

(112) Note that the iteration of the loop in FIG. **8** does not necessarily need to be executed serially. For example, where an indirect prefetch circuit already knows the address of the most recent pointer dereference “p” as well as the detected stride “d”, and the value k is a determined value or hardware-defined constant that represents the maximum number of memory locations to prefetch. Hence the memory locations p.sub.i to prefetch can be computed in parallel as (where the is a multiplication):

$$\begin{aligned} p_1 &= p + 1 \cdot \text{Math. } d \\ p_2 &= p + 2 \cdot \text{Math. } d \\ &\vdots \\ p_k &= p + K \cdot \text{Math. } d \end{aligned}$$

(114) In certain examples, each p.sub.i (from 1 to “K”) is used to prefetch memory only if the location p.sub.i contains a capability. For example, in the certain capability architectures, p.sub.i would require an associated validity tag (e.g., being set to indicate valid and not invalid).

(115) FIG. **9** is a flow diagram illustrating operations **900** of yet another method for prefetching based on a capability according to examples of the disclosure. Some or all of the operations **900** (or other processes described herein, or variations, and/or combinations thereof) are performed under the control of a prefetch circuit and/or capability management circuit as implemented herein and/or one or more computer systems configured with executable instructions and are implemented as code (e.g., executable instructions, one or more computer programs, or one or more applications) executing collectively on one or more processors, by hardware or combinations thereof. The code is stored on a computer-readable storage medium, for example, in the form of a computer program comprising instructions executable by one or more processors. The computer-readable storage medium is non-transitory. In some examples, one or more (or all) of the operations **900** are performed by prefetch circuit **120** and/or capability management circuit **108** (e.g., and execution circuit) of the other figures.

(116) The operations **900** include, at block **902**, executing, by an execution circuit of a processor, an instruction that generates a memory access request for an element in (e.g., from) memory via a first capability. The operations **900** further include, at block **904**, checking, by a capability management circuit of the processor, the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access. The operations **900** further include, at block **906**, prefetching, by a prefetch circuit of the processor, an additional element of the first object from the memory based on the first capability (e.g., where the additional element is a capability protecting a different object

than the object protected by the first capability) (e.g., the additional element having a different address than an address indicated by an address field of the first capability). The operations **900** further include, at block **908**, determining, by the processor, if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access. The operations **900** further include, at block **910**, prefetching, by the prefetch circuit, the second element from the memory to a cache of the processor based on the additional element being the second capability.

(117) The following discusses examples of prefetching for the three categories of (e.g., contiguous) data structures discussed in reference to FIG. 4 to improve prefetching accuracy.

(118) Category 1. In certain examples, an access to a single data element (e.g., single machine element) that is not part of a larger data structure does not trigger (e.g., proximal) prefetching. To understand why this can be a problem, consider register refills (e.g., from stack spill slots) (e.g., where a “spill” occurs when there are not enough registers for the variables used by a program). For example, suppose that a program returns from a procedure call and needs to refill some registers that were overwritten or deleted (e.g., clobbered) by the procedure. Refilling the registers from the stack may often happen linearly (e.g., with monotonically increasing and/or decreasing addresses) and at a fixed stride. A hardware prefetcher could easily misinterpret this behavior as an array access, and then proceed to aggressively prefetch stack memory that will not be used, and thus create unnecessary memory traffic and potentially evict hot data from the processor (e.g., CPU's) cache(s). Examples herein enhance prefetching behavior in one of two example ways, depending on how capabilities are used to access the stack: (1) the first example way: if each spill slot is assigned its own capability, then according to the operations discussed herein (e.g., in FIG. 7), each spill slot would receive its own entry in the data structure for capability based prefetching (e.g., data structure **134**), and hence an access to one spill slot followed by an access to another would not trigger the prefetcher, and thus false positives would be reduced, and (2) the second example way: if the current stack frame (for example, for the spill slots allocated on the current stack frame, e.g., where the lower bound of the stack pointer capability corresponds to the stack frame's top, and the upper bound corresponds to the stack frame's bottom) is assigned its own capability (e.g., this may be ideal to prevent the return address from being inadvertently or maliciously overwritten), then prefetching is triggered by consecutive register fills (e.g., but this prefetching would be constrained to the current stack frame where future accesses are most likely, given the principle of locality), and thus false positives are reduced and true positives are increased.

(119) Category 2. In certain examples, prefetching is (e.g., often) conducted at cache-line granularity (e.g., on many current architectures this is 64 bytes, i.e., 512 bits). In certain examples, if the heterogeneous structure is no larger than the granularity of prefetching, then the operations discussed herein (e.g., in FIG. 7) will prevent sporadic prefetching from being triggered by multiple accesses within the heterogeneous structure, and/or if the heterogeneous structure is larger than the granularity of prefetching, then prefetching is triggered within the bounds of the structure, e.g., such the operations discussed herein (e.g., in FIG. 7) will not only trigger this prefetching behavior, but it will do so aggressively, e.g., after only a threshold number of (e.g., two, three, etc.) accesses within the structure.

(120) Category 3. In certain examples, for an array of elements confined by the bounds of a capability, the operations discussed herein (e.g., in FIG. 7) will begin to aggressively prefetch memory as soon as the array is accessed a threshold number of times (e.g., twice, three times, etc.). Note that this threshold may be sooner than an array prefetcher that must rely on several accesses to establish a likely array iteration pattern. Moreover, the prefetching according to certain examples herein will not exceed the bounds of the array, as determined by the capability.

(121) Exemplary architectures, systems, etc. that the above may be used in are detailed below. Exemplary instruction formats for (e.g., capability) instructions are detailed below.

(122) At least some examples of the disclosed technologies can be described in view of the following examples: Example 1. An apparatus comprising: an execution circuit to execute an instruction that generates a memory access request for an element in memory via a first capability; a capability management circuit to check the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; a cache; and a prefetch circuit to: prefetch an additional element of the first object from the memory based on the first capability, determine if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access, and prefetch the second element from the memory to the cache based on the additional element being the second capability. Example 2. The apparatus of example 1, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory to the cache based on the first capability. Example 3. The apparatus of any one of examples 1-2, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability. Example 4. The apparatus of example 3, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability. Example 5. The apparatus of any one of examples 3-4, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits are to prefetch the second element from the memory to the cache in parallel with the prefetch of the one or more further elements of the second object from the memory to the cache. Example 6. The apparatus of any one of examples 1-5, wherein the prefetch circuit is to: prefetch a second additional element of the first object from the memory based on the first capability, determine if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third object to which the third capability authorizes access, and prefetch the third element from the memory to the cache based on the second additional element being the third capability. Example 7. The apparatus of example 6, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the third object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the third capability. Example 8. The apparatus of any one of examples 1-7, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability. Example 9. The apparatus of any one of examples 1-8, wherein the first object of the first capability and the second object of the second capability are a same object (or are different objects). Example 10. A method comprising: executing, by an execution circuit of a processor, an instruction that generates a memory access request for an element in memory via a first capability; checking, by a capability management circuit of the processor, the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; prefetching, by a prefetch circuit of the processor, an additional element of the first object from the memory based on the first capability; determining, by the processor, if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access; and prefetching, by the prefetch circuit, the second element from the memory to a cache of the processor based on the additional element being the second capability. Example 11. The method of example 10, wherein the prefetching of the additional element comprises

prefetching the additional element of the first object from the memory to the cache based on the first capability. Example 12. The method of any one of examples 10-11, further comprising prefetching, by the prefetch circuit, one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability. Example 13. The method of example 12, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability. Example 14. The method of any one of examples 12-13, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits perform the prefetching of the second element from the memory to the cache in parallel with the prefetching of the one or more further elements of the second object from the memory to the cache. Example 15. The method of any one of examples 10-14, further comprising: prefetching, by the prefetch circuit, a second additional element of the first object from the memory based on the first capability; determining, by the processor, if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third object to which the third capability authorizes access; and prefetching, by the prefetch circuit, the third element from the memory to the cache based on the second additional element being the third capability. Example 16. The method of example 15, further comprising prefetching, by the prefetch circuit, one or more further elements of the third object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the third capability. Example 17. The method of any one of examples 10-16, wherein the prefetching of the additional element of the first object from the memory is based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability. Example 18. A system comprising: a memory; a processor comprising: an execution circuit to execute an instruction that generates a memory access request for an element in memory via a first capability, and a cache separate from the memory; a capability management circuit to check the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; and a prefetch circuit to: prefetch an additional element of the first object from the memory based on the first capability, determine if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access, and prefetch the second element from the memory to the cache based on the additional element being the second capability. Example 19. The system of example 18, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory to the cache based on the first capability. Example 20. The system of any one of examples 18-19, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability. Example 21. The system of example 20, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability. Example 22. The system of any one of examples 20-21, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits are to prefetch the second element from the memory to the cache in parallel with the prefetch of the one or more further elements of the second object from the memory to the cache. Example 23. The system of any one of examples 18-22, wherein the prefetch circuit is to: prefetch a second additional element of the first object from the memory based on the first capability, determine if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third

object to which the third capability authorizes access, and prefetch the third element from the memory to the cache based on the second additional element being the third capability. Example 24. The system of example 23, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the third object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the third capability. Example 25. The system of any one of examples 18-24, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability. (123) In yet another example, an apparatus comprises a data storage device that stores code that when executed by a hardware processor causes the hardware processor to perform any method disclosed herein. An apparatus may be as described in the detailed description. A method may be as described in the detailed description.

(124) An instruction set may include one or more instruction formats. A given instruction format may define various fields (e.g., number of bits, location of bits) to specify, among other things, the operation to be performed (e.g., opcode) and the operand(s) on which that operation is to be performed and/or other data field(s) (e.g., mask). Some instruction formats are further broken down though the definition of instruction templates (or subformats). For example, the instruction templates of a given instruction format may be defined to have different subsets of the instruction format's fields (the included fields are typically in the same order, but at least some have different bit positions because there are less fields included) and/or defined to have a given field interpreted differently. Thus, each instruction of an ISA is expressed using a given instruction format (and, if defined, in a given one of the instruction templates of that instruction format) and includes fields for specifying the operation and the operands. For example, an exemplary ADD instruction has a specific opcode and an instruction format that includes an opcode field to specify that opcode and operand fields to select operands (source1/destination and source2); and an occurrence of this ADD instruction in an instruction stream will have specific contents in the operand fields that select specific operands. A set of SIMD extensions referred to as the Advanced Vector Extensions (AVX) (AVX1 and AVX2) and using the Vector Extensions (VEX) coding scheme has been released and/or published (e.g., see Intel® 64 and IA-32 Architectures Software Developer's Manual, November 2018; and see Intel® Architecture Instruction Set Extensions Programming Reference, October 2018).

(125) Exemplary Instruction Formats

(126) Examples of the instruction(s) described herein may be embodied in different formats. Additionally, exemplary systems, architectures, and pipelines are detailed below. Examples of the instruction(s) may be executed on such systems, architectures, and pipelines, but are not limited to those detailed.

(127) Generic Vector Friendly Instruction Format

(128) A vector friendly instruction format is an instruction format that is suited for vector instructions (e.g., there are certain fields specific to vector operations). While examples are described in which both vector and scalar operations are supported through the vector friendly instruction format, alternative examples use only vector operations the vector friendly instruction format.

(129) FIGS. 10A-10B are block diagrams illustrating a generic vector friendly instruction format and instruction templates thereof according to examples of the disclosure. FIG. 10A is a block diagram illustrating a generic vector friendly instruction format and class A instruction templates thereof according to examples of the disclosure; while FIG. 10B is a block diagram illustrating the generic vector friendly instruction format and class B instruction templates thereof according to examples of the disclosure. Specifically, a generic vector friendly instruction format **1000** for which are defined class A and class B instruction templates, both of which include no memory access **1005** instruction templates and memory access **1020** instruction templates. The term generic

in the context of the vector friendly instruction format refers to the instruction format not being tied to any specific instruction set.

(130) While examples of the disclosure will be described in which the vector friendly instruction format supports the following: a 64 byte vector operand length (or size) with 32 bit (4 byte) or 64 bit (8 byte) data element widths (or sizes) (and thus, a 64 byte vector consists of either 16 doubleword-size elements or alternatively, 8 quadword-size elements); a 64 byte vector operand length (or size) with 16 bit (2 byte) or 8 bit (1 byte) data element widths (or sizes); a 32 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); and a 16 byte vector operand length (or size) with 32 bit (4 byte), 64 bit (8 byte), 16 bit (2 byte), or 8 bit (1 byte) data element widths (or sizes); alternative examples may support more, less and/or different vector operand sizes (e.g., 256 byte vector operands) with more, less, or different data element widths (e.g., 128 bit (16 byte) data element widths).

(131) The class A instruction templates in FIG. 10A include: 1) within the no memory access **1005** instruction templates there is shown a no memory access, full round control type operation **1010** instruction template and a no memory access, data transform type operation **1015** instruction template; and 2) within the memory access **1020** instruction templates there is shown a memory access, temporal **1025** instruction template and a memory access, non-temporal **1030** instruction template. The class B instruction templates in FIG. 10B include: 1) within the no memory access **1005** instruction templates there is shown a no memory access, write mask control, partial round control type operation **1012** instruction template and a no memory access, write mask control, vsize type operation **1017** instruction template; and 2) within the memory access **1020** instruction templates there is shown a memory access, write mask control **1027** instruction template.

(132) The generic vector friendly instruction format **1000** includes the following fields listed below in the order illustrated in FIGS. 10A-10B.

(133) Format field **1040**—a specific value (an instruction format identifier value) in this field uniquely identifies the vector friendly instruction format, and thus occurrences of instructions in the vector friendly instruction format in instruction streams. As such, this field is optional in the sense that it is not needed for an instruction set that has only the generic vector friendly instruction format.

(134) Base operation field **1042**—its content distinguishes different base operations.

(135) Register index field **1044**—its content, directly or through address generation, specifies the locations of the source and destination operands, be they in registers or in memory. These include a sufficient number of bits to select N registers from a P×Q (e.g., 32×512, 16×128, 32×1024, 64×1024) register file. While in one example N may be up to three sources and one destination register, alternative examples may support more or less sources and destination registers (e.g., may support up to two sources where one of these sources also acts as the destination, may support up to three sources where one of these sources also acts as the destination, may support up to two sources and one destination).

(136) Modifier field **1046**—its content distinguishes occurrences of instructions in the generic vector instruction format that specify memory access from those that do not; that is, between no memory access **1005** instruction templates and memory access **1020** instruction templates. Memory access operations read and/or write to the memory hierarchy (in some cases specifying the source and/or destination addresses using values in registers), while non-memory access operations do not (e.g., the source and destinations are registers). While in one example this field also selects between three different ways to perform memory address calculations, alternative examples may support more, less, or different ways to perform memory address calculations.

(137) Augmentation operation field **1050**—its content distinguishes which one of a variety of different operations to be performed in addition to the base operation. This field is context specific. In one example of the disclosure, this field is divided into a class field **1068**, an alpha field **1052**, and a beta field **1054**. The augmentation operation field **1050** allows common groups of operations

to be performed in a single instruction rather than 2, 3, or 4 instructions.

(138) Scale field **1060**—its content allows for the scaling of the index field's content for memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base}$).

(139) Displacement Field **1062A**—its content is used as part of memory address generation (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{displacement}$).

(140) Displacement Factor Field **1062B** (note that the juxtaposition of displacement field **1062A** directly over displacement factor field **1062B** indicates one or the other is used)—its content is used as part of address generation; it specifies a displacement factor that is to be scaled by the size of a memory access (N)—where N is the number of bytes in the memory access (e.g., for address generation that uses $2.\text{sup.scale} \times \text{index} + \text{base} + \text{scaled displacement}$). Redundant low-order bits are ignored and hence, the displacement factor field's content is multiplied by the memory operands total size (N) in order to generate the final displacement to be used in calculating an effective address. The value of N is determined by the processor hardware at runtime based on the full opcode field **1074** (described later herein) and the data manipulation field **1054C**. The displacement field **1062A** and the displacement factor field **1062B** are optional in the sense that they are not used for the no memory access **1005** instruction templates and/or different examples may implement only one or none of the two.

(141) Data element width field **1064**—its content distinguishes which one of a number of data element widths is to be used (in some examples for all instructions; in other examples for only some of the instructions). This field is optional in the sense that it is not needed if only one data element width is supported and/or data element widths are supported using some aspect of the opcodes.

(142) Write mask field **1070**—its content controls, on a per data element position basis, whether that data element position in the destination vector operand reflects the result of the base operation and augmentation operation. Class A instruction templates support merging-writemasking, while class B instruction templates support both merging- and zeroing-writemasking. When merging, vector masks allow any set of elements in the destination to be protected from updates during the execution of any operation (specified by the base operation and the augmentation operation); in other one example, preserving the old value of each element of the destination where the corresponding mask bit has a 0. In contrast, when zeroing vector masks allow any set of elements in the destination to be zeroed during the execution of any operation (specified by the base operation and the augmentation operation); in one example, an element of the destination is set to 0 when the corresponding mask bit has a 0 value. A subset of this functionality is the ability to control the vector length of the operation being performed (that is, the span of elements being modified, from the first to the last one); however, it is not necessary that the elements that are modified be consecutive. Thus, the write mask field **1070** allows for partial vector operations, including loads, stores, arithmetic, logical, etc. While examples of the disclosure are described in which the write mask field's **1070** content selects one of a number of write mask registers that contains the write mask to be used (and thus the write mask field's **1070** content indirectly identifies that masking to be performed), alternative examples instead or additional allow the mask write field's **1070** content to directly specify the masking to be performed.

(143) Immediate field **1072**—its content allows for the specification of an immediate. This field is optional in the sense that it is not present in an implementation of the generic vector friendly format that does not support immediate and it is not present in instructions that do not use an immediate.

(144) Class field **1068**—its content distinguishes between different classes of instructions. With reference to FIGS. **10A-B**, the contents of this field select between class A and class B instructions. In FIGS. **10A-B**, rounded corner squares are used to indicate a specific value is present in a field (e.g., class A **1068A** and class B **1068B** for the class field **1068** respectively in FIGS. **10A-B**).

(145) Instruction Templates of Class A

(146) In the case of the non-memory access **1005** instruction templates of class A, the alpha field

1052 is interpreted as an RS field **1052A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **1052A.1** and data transform **1052A.2** are respectively specified for the no memory access, round type operation **1010** and the no memory access, data transform type operation **1015** instruction templates), while the beta field **1054** distinguishes which of the operations of the specified type is to be performed. In the no memory access **1005** instruction templates, the scale field **1060**, the displacement field **1062A**, and the displacement scale field **1062B** are not present.

(147) No-Memory Access Instruction Templates—Full Round Control Type Operation

(148) In the no memory access full round control type operation **1010** instruction template, the beta field **1054** is interpreted as a round control field **1054A**, whose content(s) provide static rounding. While in the described examples of the disclosure the round control field **1054A** includes a suppress all floating point exceptions (SAE) field **1056** and a round operation control field **1058**, alternative examples may support may encode both these concepts into the same field or only have one or the other of these concepts/fields (e.g., may have only the round operation control field **1058**).

(149) SAE field **1056**—its content distinguishes whether or not to disable the exception event reporting; when the SAE field's **1056** content indicates suppression is enabled, a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler.

(150) Round operation control field **1058**—its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **1058** allows for the changing of the rounding mode on a per instruction basis. In one example of the disclosure where a processor includes a control register for specifying rounding modes, the round operation control field's **1050** content overrides that register value.

(151) No Memory Access Instruction Templates—Data Transform Type Operation

(152) In the no memory access data transform type operation **1015** instruction template, the beta field **1054** is interpreted as a data transform field **1054B**, whose content distinguishes which one of a number of data transforms is to be performed (e.g., no data transform, swizzle, broadcast).

(153) In the case of a memory access **1020** instruction template of class A, the alpha field **1052** is interpreted as an eviction hint field **1052B**, whose content distinguishes which one of the eviction hints is to be used (in FIG. 10A, temporal **1052B.1** and non-temporal **1052B.2** are respectively specified for the memory access, temporal **1025** instruction template and the memory access, non-temporal **1030** instruction template), while the beta field **1054** is interpreted as a data manipulation field **1054C**, whose content distinguishes which one of a number of data manipulation operations (also known as primitives) is to be performed (e.g., no manipulation; broadcast; up conversion of a source; and down conversion of a destination). The memory access **1020** instruction templates include the scale field **1060**, and optionally the displacement field **1062A** or the displacement scale field **1062B**.

(154) Vector memory instructions perform vector loads from and vector stores to memory, with conversion support. As with regular vector instructions, vector memory instructions transfer data from/to memory in a data element-wise fashion, with the elements that are actually transferred is dictated by the contents of the vector mask that is selected as the write mask.

(155) Memory Access Instruction Templates—Temporal

(156) Temporal data is data likely to be reused soon enough to benefit from caching. This is, however, a hint, and different processors may implement it in different ways, including ignoring the hint entirely.

(157) Memory Access Instruction Templates—Non-Temporal

(158) Non-temporal data is data unlikely to be reused soon enough to benefit from caching in the 1st-level cache and should be given priority for eviction. This is, however, a hint, and different

processors may implement it in different ways, including ignoring the hint entirely.

(159) Instruction Templates of Class B

(160) In the case of the instruction templates of class B, the alpha field **1052** is interpreted as a write mask control (Z) field **1052C**, whose content distinguishes whether the write masking controlled by the write mask field **1070** should be a merging or a zeroing.

(161) In the case of the non-memory access **1005** instruction templates of class B, part of the beta field **1054** is interpreted as an RL field **1057A**, whose content distinguishes which one of the different augmentation operation types are to be performed (e.g., round **1057A.1** and vector length (VSIZE) **1057A.2** are respectively specified for the no memory access, write mask control, partial round control type operation **1012** instruction template and the no memory access, write mask control, VSIZE type operation **1017** instruction template), while the rest of the beta field **1054** distinguishes which of the operations of the specified type is to be performed. In the no memory access **1005** instruction templates, the scale field **1060**, the displacement field **1062A**, and the displacement scale field **1062B** are not present.

(162) In the no memory access, write mask control, partial round control type operation **1010** instruction template, the rest of the beta field **1054** is interpreted as a round operation field **1059A** and exception event reporting is disabled (a given instruction does not report any kind of floating-point exception flag and does not raise any floating point exception handler).

(163) Round operation control field **1059A**—just as round operation control field **1058**, its content distinguishes which one of a group of rounding operations to perform (e.g., Round-up, Round-down, Round-towards-zero and Round-to-nearest). Thus, the round operation control field **1059A** allows for the changing of the rounding mode on a per instruction basis. In one example of the disclosure where a processor includes a control register for specifying rounding modes, the round operation control field's **1050** content overrides that register value.

(164) In the no memory access, write mask control, VSIZE type operation **1017** instruction template, the rest of the beta field **1054** is interpreted as a vector length field **1059B**, whose content distinguishes which one of a number of data vector lengths is to be performed on (e.g., 128, 256, or 512 byte).

(165) In the case of a memory access **1020** instruction template of class B, part of the beta field **1054** is interpreted as a broadcast field **1057B**, whose content distinguishes whether or not the broadcast type data manipulation operation is to be performed, while the rest of the beta field **1054** is interpreted the vector length field **1059B**. The memory access **1020** instruction templates include the scale field **1060**, and optionally the displacement field **1062A** or the displacement scale field **1062B**.

(166) With regard to the generic vector friendly instruction format **1000**, a full opcode field **1074** is shown including the format field **1040**, the base operation field **1042**, and the data element width field **1064**. While one example is shown where the full opcode field **1074** includes all of these fields, the full opcode field **1074** includes less than all of these fields in examples that do not support all of them. The full opcode field **1074** provides the operation code (opcode).

(167) The augmentation operation field **1050**, the data element width field **1064**, and the write mask field **1070** allow these features to be specified on a per instruction basis in the generic vector friendly instruction format.

(168) The combination of write mask field and data element width field create typed instructions in that they allow the mask to be applied based on different data element widths.

(169) The various instruction templates found within class A and class B are beneficial in different situations. In some examples of the disclosure, different processors or different cores within a processor may support only class A, only class B, or both classes. For instance, a high performance general purpose out-of-order core intended for general-purpose computing may support only class B, a core intended primarily for graphics and/or scientific (throughput) computing may support only class A, and a core intended for both may support both (of course, a core that has some mix of

templates and instructions from both classes but not all templates and instructions from both classes is within the purview of the disclosure). Also, a single processor may include multiple cores, all of which support the same class or in which different cores support different class. For instance, in a processor with separate graphics and general purpose cores, one of the graphics cores intended primarily for graphics and/or scientific computing may support only class A, while one or more of the general purpose cores may be high performance general purpose cores with out of order execution and register renaming intended for general-purpose computing that support only class B. Another processor that does not have a separate graphics core, may include one more general purpose in-order or out-of-order cores that support both class A and class B. Of course, features from one class may also be implement in the other class in different examples of the disclosure. Programs written in a high level language would be put (e.g., just in time compiled or statically compiled) into an variety of different executable forms, including: 1) a form having only instructions of the class(es) supported by the target processor for execution; or 2) a form having alternative routines written using different combinations of the instructions of all classes and having control flow code that selects the routines to execute based on the instructions supported by the processor which is currently executing the code.

(170) Exemplary Specific Vector Friendly Instruction Format

(171) FIG. **11** is a block diagram illustrating an exemplary specific vector friendly instruction format according to examples of the disclosure. FIG. **11** shows a specific vector friendly instruction format **1100** that is specific in the sense that it specifies the location, size, interpretation, and order of the fields, as well as values for some of those fields. The specific vector friendly instruction format **1100** may be used to extend the x86 instruction set, and thus some of the fields are similar or the same as those used in the existing x86 instruction set and extension thereof (e.g., AVX). This format remains consistent with the prefix encoding field, real opcode byte field, MOD R/M field, SIB field, displacement field, and immediate fields of the existing x86 instruction set with extensions. The fields from FIG. **10** into which the fields from FIG. **11** map are illustrated.

(172) It should be understood that, although examples of the disclosure are described with reference to the specific vector friendly instruction format **1100** in the context of the generic vector friendly instruction format **1000** for illustrative purposes, the disclosure is not limited to the specific vector friendly instruction format **1100** except where claimed. For example, the generic vector friendly instruction format **1000** contemplates a variety of possible sizes for the various fields, while the specific vector friendly instruction format **1100** is shown as having fields of specific sizes. By way of specific example, while the data element width field **1064** is illustrated as a one bit field in the specific vector friendly instruction format **1100**, the disclosure is not so limited (that is, the generic vector friendly instruction format **1000** contemplates other sizes of the data element width field **1064**).

(173) The generic vector friendly instruction format **1000** includes the following fields listed below in the order illustrated in FIG. **11A**.

(174) EVEX Prefix (Bytes 0-3) **1102**—is encoded in a four-byte form.

(175) Format Field **1040** (EVEX Byte 0, bits [7:0])—the first byte (EVEX Byte 0) is the format field **1040** and it contains 0x62 (the unique value used for distinguishing the vector friendly instruction format in one example of the disclosure).

(176) The second-fourth bytes (EVEX Bytes 1-3) include a number of bit fields providing specific capability.

(177) REX field **1105** (EVEX Byte 1, bits [7-5])—consists of a EVEX.R bit field (EVEX Byte 1, bit [7]—R), EVEX.X bit field (EVEX byte 1, bit [6]—X), and **1057BEX** byte 1, bit[5]—B). The EVEX.R, EVEX.X, and EVEX.B bit fields provide the same functionality as the corresponding VEX bit fields, and are encoded using is complement form, e.g., ZMM0 is encoded as **1111B**, ZMM15 is encoded as 0000B. Other fields of the instructions encode the lower three bits of the register indexes as is known in the art (rrr, xxx, and bbb), so that Rrrr, Xxxx, and Bbbb may be

formed by adding EVEX.R, EVEX.X, and EVEX.B.

(178) REX' field **1010**—this is the first part of the REX' field **1010** and is the EVEX.R' bit field (EVEX Byte 1, bit [4]—R') that is used to encode either the upper 16 or lower 16 of the extended 32 register set. In one example of the disclosure, this bit, along with others as indicated below, is stored in bit inverted format to distinguish (in the well-known x86 32-bit mode) from the BOUND instruction, whose real opcode byte is 62, but does not accept in the MOD R/M field (described below) the value of 11 in the MOD field; alternative examples of the disclosure do not store this and the other indicated bits below in the inverted format. A value of 1 is used to encode the lower 16 registers. In other words, R'Rrrr is formed by combining EVEX.R', EVEX.R, and the other RRR from other fields.

(179) Opcode map field **1115** (EVEX byte 1, bits [3:0]—mmmm)—its content encodes an implied leading opcode byte (0F, 0F 38, or 0F 3).

(180) Data element width field **1064** (EVEX byte 2, bit [7]—W)—is represented by the notation EVEX.W. EVEX.W is used to define the granularity (size) of the datatype (either 32-bit data elements or 64-bit data elements).

(181) EVEX.vvvv **1120** (EVEX Byte 2, bits [6:3]-vvvv)—the role of EVEX.vvvv may include the following: 1) EVEX.vvvv encodes the first source register operand, specified in inverted (1s complement) form and is valid for instructions with 2 or more source operands; 2) EVEX.vvvv encodes the destination register operand, specified in is complement form for certain vector shifts; or 3) EVEX.vvvv does not encode any operand, the field is reserved and should contain **1111b**. Thus, EVEX.vvvv field **1120** encodes the 4 low-order bits of the first source register specifier stored in inverted (1s complement) form. Depending on the instruction, an extra different EVEX bit field is used to extend the specifier size to 32 registers.

(182) EVEX.U **1068** Class field (EVEX byte 2, bit [2]-U)—If EVEX.U=0, it indicates class A or EVEX.U0; if EVEX.U=1, it indicates class B or EVEX.U1.

(183) Prefix encoding field **1125** (EVEX byte 2, bits [1:0]-pp)—provides additional bits for the base operation field. In addition to providing support for the legacy SSE instructions in the EVEX prefix format, this also has the benefit of compacting the SIMD prefix (rather than requiring a byte to express the SIMD prefix, the EVEX prefix requires only 2 bits). In one example, to support legacy SSE instructions that use a SIMD prefix (66H, F2H, F3H) in both the legacy format and in the EVEX prefix format, these legacy SIMD prefixes are encoded into the SIMD prefix encoding field; and at runtime are expanded into the legacy SIMD prefix prior to being provided to the decoder's PLA (so the PLA can execute both the legacy and EVEX format of these legacy instructions without modification). Although newer instructions could use the EVEX prefix encoding field's content directly as an opcode extension, certain examples expand in a similar fashion for consistency but allow for different meanings to be specified by these legacy SIMD prefixes. An alternative example may redesign the PLA to support the 2 bit SIMD prefix encodings, and thus not require the expansion.

(184) Alpha field **1052** (EVEX byte 3, bit [7]—EH; also known as EVEX.EH, EVEX.rs, EVEX.RL, EVEX.write mask control, and EVEX.N; also illustrated with α)—as previously described, this field is context specific.

(185) Beta field **1054** (EVEX byte 3, bits [6:4]-SSS, also known as EVEX.s.sub.2-0, EVEX.r.sub.2-0, EVEX.rr1, EVEX.LL0, EVEX.LLB; also illustrated with $\beta\beta\beta$)—as previously described, this field is context specific.

(186) REX' field **1010**—this is the remainder of the REX' field and is the EVEX.V' bit field (EVEX Byte 3, bit [3]—V') that may be used to encode either the upper 16 or lower 16 of the extended 32 register set. This bit is stored in bit inverted format. A value of 1 is used to encode the lower 16 registers. In other words, V'VVVV is formed by combining EVEX.V', EVEX.vvvv.

(187) Write mask field **1070** (EVEX byte 3, bits [2:0]-kkk)—its content specifies the index of a register in the write mask registers as previously described. In one example of the disclosure, the

specific value EVEX.kkk=000 has a special behavior implying no write mask is used for the particular instruction (this may be implemented in a variety of ways including the use of a write mask hardwired to all ones or hardware that bypasses the masking hardware).

(188) Real Opcode Field **1130** (Byte 4) is also known as the opcode byte. Part of the opcode is specified in this field.

(189) MOD R/M Field **1140** (Byte 5) includes MOD field **1142**, Reg field **1144**, and R/M field **1146**. As previously described, the MOD field's **1142** content distinguishes between memory access and non-memory access operations. The role of Reg field **1144** can be summarized to two situations: encoding either the destination register operand or a source register operand, or be treated as an opcode extension and not used to encode any instruction operand. The role of R/M field **1146** may include the following: encoding the instruction operand that references a memory address, or encoding either the destination register operand or a source register operand.

(190) Scale, Index, Base (SIB) Byte (Byte 6)—As previously described, the scale field's **1050** content is used for memory address generation. SIB.xxx **1154** and SIB.bbb **1156**—the contents of these fields have been previously referred to with regard to the register indexes Xxxx and Bbbb.

(191) Displacement field **1062A** (Bytes 7-10)—when MOD field **1142** contains 10, bytes 7-10 are the displacement field **1062A**, and it works the same as the legacy 32-bit displacement (disp32) and works at byte granularity.

(192) Displacement factor field **1062B** (Byte 7)—when MOD field **1142** contains 01, byte 7 is the displacement factor field **1062B**. The location of this field is that same as that of the legacy x86 instruction set 8-bit displacement (disp8), which works at byte granularity. Since disp8 is sign extended, it can only address between -128 and 127 bytes offsets; in terms of 64 byte cache lines, disp8 uses 8 bits that can be set to only four really useful values -128, -64, 0, and 64; since a greater range is often needed, disp32 is used; however, disp32 requires 4 bytes. In contrast to disp8 and disp32, the displacement factor field **1062B** is a reinterpretation of disp8; when using displacement factor field **1062B**, the actual displacement is determined by the content of the displacement factor field multiplied by the size of the memory operand access (N). This type of displacement is referred to as $\text{disp8} * N$. This reduces the average instruction length (a single byte of used for the displacement but with a much greater range). Such compressed displacement is based on the assumption that the effective displacement is multiple of the granularity of the memory access, and hence, the redundant low-order bits of the address offset do not need to be encoded. In other words, the displacement factor field **1062B** substitutes the legacy x86 instruction set 8-bit displacement. Thus, the displacement factor field **1062B** is encoded the same way as an x86 instruction set 8-bit displacement (so no changes in the ModRM/SIB encoding rules) with the only exception that disp8 is overloaded to $\text{disp8} * N$. In other words, there are no changes in the encoding rules or encoding lengths but only in the interpretation of the displacement value by hardware (which needs to scale the displacement by the size of the memory operand to obtain a byte-wise address offset). Immediate field **1072** operates as previously described.

(193) Full Opcode Field

(194) FIG. **11B** is a block diagram illustrating the fields of the specific vector friendly instruction format **1100** that make up the full opcode field **1074** according to one example of the disclosure. Specifically, the full opcode field **1074** includes the format field **1040**, the base operation field **1042**, and the data element width (W) field **1064**. The base operation field **1042** includes the prefix encoding field **1125**, the opcode map field **1115**, and the real opcode field **1130**.

(195) Register Index Field

(196) FIG. **11C** is a block diagram illustrating the fields of the specific vector friendly instruction format **1100** that make up the register index field **1044** according to one example of the disclosure. Specifically, the register index field **1044** includes the REX field **1105**, the REX' field **1110**, the MODR/M.reg field **1144**, the MODR/M.r/m field **1146**, the VVVV field **1120**, xxx field **1154**, and the bbb field **1156**.

(197) Augmentation Operation Field

(198) FIG. 11D is a block diagram illustrating the fields of the specific vector friendly instruction format **1100** that make up the augmentation operation field **1050** according to one example of the disclosure. When the class (U) field **1068** contains 0, it signifies EVEX.U0 (class A **1068A**); when it contains 1, it signifies EVEX.U1 (class B **1068B**). When U=0 and the MOD field **1142** contains 11 (signifying a no memory access operation), the alpha field **1052** (EVEX byte 3, bit [7]—EH) is interpreted as the rs field **1052A**. When the rs field **1052A** contains a 1 (round **1052A.1**), the beta field **1054** (EVEX byte 3, bits [6:4]—SSS) is interpreted as the round control field **1054A**. The round control field **1054A** includes a one bit SAE field **1056** and a two bit round operation field **1058**. When the rs field **1052A** contains a 0 (data transform **1052A.2**), the beta field **1054** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data transform field **1054B**. When U=0 and the MOD field **1142** contains 00, 01, or 10 (signifying a memory access operation), the alpha field **1052** (EVEX byte 3, bit [7]—EH) is interpreted as the eviction hint (EH) field **1052B** and the beta field **1054** (EVEX byte 3, bits [6:4]—SSS) is interpreted as a three bit data manipulation field **1054C**.

(199) When U=1, the alpha field **1052** (EVEX byte 3, bit [7]—EH) is interpreted as the write mask control (Z) field **1052C**. When U=1 and the MOD field **1142** contains 11 (signifying a no memory access operation), part of the beta field **1054** (EVEX byte 3, bit [4]—S.sub.0) is interpreted as the RL field **1057A**; when it contains a 1 (round **1057A.1**) the rest of the beta field **1054** (EVEX byte 3, bit [6-5]—S.sub.2-1) is interpreted as the round operation field **1059A**, while when the RL field **1057A** contains a 0 (VSIZE **1057.A2**) the rest of the beta field **1054** (EVEX byte 3, bit [6-5]—S.sub.2-1) is interpreted as the vector length field **1059B** (EVEX byte 3, bit [6-5]—L.sub.1-0). When U=1 and the MOD field **1142** contains 00, 01, or 10 (signifying a memory access operation), the beta field **1054** (EVEX byte 3, bits [6:4]—SSS) is interpreted as the vector length field **1059B** (EVEX byte 3, bit [6-5]—L.sub.1-0) and the broadcast field **1057B** (EVEX byte 3, bit [4]—B).

(200) Exemplary Register Architecture

(201) FIG. 12 is a block diagram of a register architecture **1200** according to one example of the disclosure. In the example illustrated, there are 32 vector registers **1210** that are 512 bits wide; these registers are referenced as zmm0 through zmm31. The lower order 256 bits of the lower 16 zmm registers are overlaid on registers ymm0-16. The lower order 128 bits of the lower 16 zmm registers (the lower order 128 bits of the ymm registers) are overlaid on registers xmm0-15. The specific vector friendly instruction format **1100** operates on these overlaid register file as illustrated in the below tables.

(202) TABLE-US-00001 Adjustable Vector Length Class Operations Registers Instruction Templates that A (FIG. 10A; 1010, 1015, zmm registers (the vector length do not include the vector U = 0) 1025, 1030 is 64 byte) length field 1059B B (FIG. 10B; 1012 zmm registers (the vector length U = 1) is 64 byte) Instruction templates that B (FIG. 10B; 1017, 1027 zmm, ymm, or xmm registers (the do include the vector U = 1) vector length is 64 byte, 32 byte, length field 1059B or 16 byte) depending on the vector length field 1059B

(203) In other words, the vector length field **1059B** selects between a maximum length and one or more other shorter lengths, where each such shorter length is half the length of the preceding length; and instructions templates without the vector length field **1059B** operate on the maximum vector length. Further, in one example, the class B instruction templates of the specific vector friendly instruction format **1100** operate on packed or scalar single/double-precision floating point data and packed or scalar integer data. Scalar operations are operations performed on the lowest order data element position in a zmm/ymm/xmm register; the higher order data element positions are either left the same as they were prior to the instruction or zeroed depending on the example.

(204) Write mask registers **1215**—in the example illustrated, there are 8 write mask registers (k0 through k7), each 64 bits in size. In an alternate example, the write mask registers **1215** are 16 bits in size. As previously described, in one example of the disclosure, the vector mask register k0

cannot be used as a write mask; when the encoding that would normally indicate k0 is used for a write mask, it selects a hardwired write mask of 0xFFFF, effectively disabling write masking for that instruction.

(205) General-purpose registers **1225**—in the example illustrated, there are sixteen 64-bit general-purpose registers that are used along with the existing x86 addressing modes to address memory operands. These registers are referenced by the names RAX, RBX, RCX, RDX, RBP, RSI, RDI, RSP, and R8 through R15.

(206) Scalar floating point stack register file (x87 stack) **1245**, on which is aliased the MMX packed integer flat register file **1250**—in the example illustrated, the x87 stack is an eight-element stack used to perform scalar floating-point operations on 32/64/80-bit floating point data using the x87 instruction set extension; while the MMX registers are used to perform operations on 64-bit packed integer data, as well as to hold operands for some operations performed between the MMX and XMM registers.

(207) Alternative examples of the disclosure may use wider or narrower registers. Additionally, alternative examples of the disclosure may use more, less, or different register files and registers.

(208) Exemplary Core Architectures, Processors, and Computer Architectures

(209) Processor cores may be implemented in different ways, for different purposes, and in different processors. For instance, implementations of such cores may include: 1) a general purpose in-order core intended for general-purpose computing; 2) a high performance general purpose out-of-order core intended for general-purpose computing; 3) a special purpose core intended primarily for graphics and/or scientific (throughput) computing. Implementations of different processors may include: 1) a CPU including one or more general purpose in-order cores intended for general-purpose computing and/or one or more general purpose out-of-order cores intended for general-purpose computing; and 2) a coprocessor including one or more special purpose cores intended primarily for graphics and/or scientific (throughput). Such different processors lead to different computer system architectures, which may include: 1) the coprocessor on a separate chip from the CPU; 2) the coprocessor on a separate die in the same package as a CPU; 3) the coprocessor on the same die as a CPU (in which case, such a coprocessor is sometimes referred to as special purpose logic, such as integrated graphics and/or scientific (throughput) logic, or as special purpose cores); and 4) a system on a chip that may include on the same die the described CPU (sometimes referred to as the application core(s) or application processor(s)), the above described coprocessor, and additional functionality. Exemplary core architectures are described next, followed by descriptions of exemplary processors and computer architectures.

(210) Exemplary Core Architectures

(211) In-Order and Out-of-Order Core Block Diagram

(212) FIG. **13A** is a block diagram illustrating both an exemplary in-order pipeline and an exemplary register renaming, out-of-order issue/execution pipeline according to examples of the disclosure. FIG. **13B** is a block diagram illustrating both an exemplary example of an in-order architecture core and an exemplary register renaming, out-of-order issue/execution architecture core to be included in a processor according to examples of the disclosure. The solid lined boxes in FIGS. **13A-B** illustrate the in-order pipeline and in-order core, while the optional addition of the dashed lined boxes illustrates the register renaming, out-of-order issue/execution pipeline and core. Given that the in-order aspect is a subset of the out-of-order aspect, the out-of-order aspect will be described.

(213) In FIG. **13A**, a processor pipeline **1300** includes a fetch stage **1302**, a length decode stage **1304**, a decode stage **1306**, an allocation stage **1308**, a renaming stage **1310**, a scheduling (also known as a dispatch or issue) stage **1312**, a register read/memory read stage **1314**, an execute stage **1316**, a write back/memory write stage **1318**, an exception handling stage **1322**, and a commit stage **1324**.

(214) FIG. **13B** shows processor core **1390** including a front end unit **1330** coupled to an execution

engine unit **1350**, and both are coupled to a memory unit **1370**. The core **1390** may be a reduced instruction set computing (RISC) core, a complex instruction set computing (CISC) core, a very long instruction word (VLIW) core, or a hybrid or alternative core type. As yet another option, the core **1390** may be a special-purpose core, such as, for example, a network or communication core, compression engine, coprocessor core, general purpose computing graphics processing unit (GPGPU) core, graphics core, or the like.

(215) The front end unit **1330** includes a branch prediction unit **1332** coupled to an instruction cache unit **1334**, which is coupled to an instruction translation lookaside buffer (TLB) **1336**, which is coupled to an instruction fetch unit **1338**, which is coupled to a decode unit **1340**. The decode unit **1340** (or decoder or decoder unit) may decode instructions (e.g., macro-instructions), and generate as an output one or more micro-operations, micro-code entry points, micro-instructions, other instructions, or other control signals, which are decoded from, or which otherwise reflect, or are derived from, the original instructions. The decode unit **1340** may be implemented using various different mechanisms. Examples of suitable mechanisms include, but are not limited to, look-up tables, hardware implementations, programmable logic arrays (PLAs), microcode read only memories (ROMs), etc. In one example, the core **1390** includes a microcode ROM or other medium that stores microcode for certain macro-instructions (e.g., in decode unit **1340** or otherwise within the front end unit **1330**). The decode unit **1340** is coupled to a rename/allocator unit **1352** in the execution engine unit **1350**.

(216) The execution engine unit **1350** includes the rename/allocator unit **1352** coupled to a retirement unit **1354** and a set of one or more scheduler unit(s) **1356**. The scheduler unit(s) **1356** represents any number of different schedulers, including reservations stations, central instruction window, etc. The scheduler unit(s) **1356** is coupled to the physical register file(s) unit(s) **1358**. Each of the physical register file(s) units **1358** represents one or more physical register files, different ones of which store one or more different data types, such as scalar integer, scalar floating point, packed integer, packed floating point, vector integer, vector floating point, status (e.g., an instruction pointer that is the address of the next instruction to be executed), etc. In one example, the physical register file(s) unit **1358** comprises a vector registers unit, a write mask registers unit, and a scalar registers unit. These register units may provide architectural vector registers, vector mask registers, and general purpose registers. The physical register file(s) unit(s) **1358** is overlapped by the retirement unit **1354** to illustrate various ways in which register renaming and out-of-order execution may be implemented (e.g., using a reorder buffer(s) and a retirement register file(s); using a future file(s), a history buffer(s), and a retirement register file(s); using a register maps and a pool of registers; etc.). The retirement unit **1354** and the physical register file(s) unit(s) **1358** are coupled to the execution cluster(s) **1360**. The execution cluster(s) **1360** includes a set of one or more execution units **1362** and a set of one or more memory access units **1364**. The execution units **1362** may perform various operations (e.g., shifts, addition, subtraction, multiplication) and on various types of data (e.g., scalar floating point, packed integer, packed floating point, vector integer, vector floating point). While some examples may include a number of execution units dedicated to specific functions or sets of functions, other examples may include only one execution unit or multiple execution units that all perform all functions. The scheduler unit(s) **1356**, physical register file(s) unit(s) **1358**, and execution cluster(s) **1360** are shown as being possibly plural because certain examples create separate pipelines for certain types of data/operations (e.g., a scalar integer pipeline, a scalar floating point/packed integer/packed floating point/vector integer/vector floating point pipeline, and/or a memory access pipeline that each have their own scheduler unit, physical register file(s) unit, and/or execution cluster—and in the case of a separate memory access pipeline, certain examples are implemented in which only the execution cluster of this pipeline has the memory access unit(s) **1364**). It should also be understood that where separate pipelines are used, one or more of these pipelines may be out-of-order issue/execution and the rest in-order.

(217) The set of memory access units **1364** is coupled to the memory unit **1370**, which includes a data TLB unit **1372** coupled to a data cache unit **1374** coupled to a level 2 (L2) cache unit **1376**. In one exemplary example, the memory access units **1364** may include a load unit, a store address unit, and a store data unit, each of which is coupled to the data TLB unit **1372** in the memory unit **1370**. The instruction cache unit **1334** is further coupled to a level 2 (L2) cache unit **1376** in the memory unit **1370**. The L2 cache unit **1376** is coupled to one or more other levels of cache and eventually to a main memory.

(218) In certain examples, a prefetch circuit **1378** is included to prefetch data, for example, to predict access addresses and bring the data for those addresses into a cache or caches (e.g., from memory **1380**).

(219) By way of example, the exemplary register renaming, out-of-order issue/execution core architecture may implement the pipeline **1300** as follows: 1) the instruction fetch **1338** performs the fetch and length decoding stages **1302** and **1304**; 2) the decode unit **1340** performs the decode stage **1306**; 3) the rename/allocator unit **1352** performs the allocation stage **1308** and renaming stage **1310**; 4) the scheduler unit(s) **1356** performs the schedule stage **1312**; 5) the physical register file(s) unit(s) **1358** and the memory unit **1370** perform the register read/memory read stage **1314**; the execution cluster **1360** perform the execute stage **1316**; 6) the memory unit **1370** and the physical register file(s) unit(s) **1358** perform the write back/memory write stage **1318**; 7) various units may be involved in the exception handling stage **1322**; and 8) the retirement unit **1354** and the physical register file(s) unit(s) **1358** perform the commit stage **1324**.

(220) The core **1390** may support one or more instructions sets (e.g., the x86 instruction set (with some extensions that have been added with newer versions); the MIPS instruction set of MIPS Technologies of Sunnyvale, CA; the ARM instruction set (with optional additional extensions such as NEON) of ARM Holdings of Sunnyvale, CA), including the instruction(s) described herein. In one example, the core **1390** includes logic to support a packed data instruction set extension (e.g., AVX1, AVX2), thereby allowing the operations used by many multimedia applications to be performed using packed data.

(221) It should be understood that the core may support multithreading (executing two or more parallel sets of operations or threads), and may do so in a variety of ways including time sliced multithreading, simultaneous multithreading (where a single physical core provides a logical core for each of the threads that physical core is simultaneously multithreading), or a combination thereof (e.g., time sliced fetching and decoding and simultaneous multithreading thereafter such as in the Intel® Hyper-Threading technology).

(222) While register renaming is described in the context of out-of-order execution, it should be understood that register renaming may be used in an in-order architecture. While the illustrated example of the processor also includes separate instruction and data cache units **1334/1374** and a shared L2 cache unit **1376**, alternative examples may have a single internal cache for both instructions and data, such as, for example, a Level 1 (L1) internal cache, or multiple levels of internal cache. In some examples, the system may include a combination of an internal cache and an external cache that is external to the core and/or the processor. Alternatively, all of the cache may be external to the core and/or the processor.

(223) Specific Exemplary In-Order Core Architecture

(224) FIGS. **14A-B** illustrate a block diagram of a more specific exemplary in-order core architecture, which core would be one of several logic blocks (including other cores of the same type and/or different types) in a chip. The logic blocks communicate through a high-bandwidth interconnect network (e.g., a ring network) with some fixed function logic, memory I/O interfaces, and other necessary I/O logic, depending on the application.

(225) FIG. **14A** is a block diagram of a single processor core, along with its connection to the on-die interconnect network **1402** and with its local subset of the Level 2 (L2) cache **1404**, according to examples of the disclosure. In one example, an instruction decode unit **1400** supports the x86

instruction set with a packed data instruction set extension. An L1 cache **1406** allows low-latency accesses to cache memory into the scalar and vector units. While in one example (to simplify the design), a scalar unit **1408** and a vector unit **1410** use separate register sets (respectively, scalar registers **1412** and vector registers **1414**) and data transferred between them is written to memory and then read back in from a level 1 (L1) cache **1406**, alternative examples of the disclosure may use a different approach (e.g., use a single register set or include a communication path that allow data to be transferred between the two register files without being written and read back).

(226) The local subset of the L2 cache **1404** is part of a global L2 cache that is divided into separate local subsets, one per processor core. Each processor core has a direct access path to its own local subset of the L2 cache **1404**. Data read by a processor core is stored in its L2 cache subset **1404** and can be accessed quickly, in parallel with other processor cores accessing their own local L2 cache subsets. Data written by a processor core is stored in its own L2 cache subset **1404** and is flushed from other subsets, if necessary. The ring network ensures coherency for shared data. The ring network is bi-directional to allow agents such as processor cores, L2 caches and other logic blocks to communicate with each other within the chip. Each ring data-path is 1012-bits wide per direction.

(227) FIG. **14B** is an expanded view of part of the processor core in FIG. **14A** according to examples of the disclosure. FIG. **14B** includes an L1 data cache **1406A** part of the L1 cache **1404**, as well as more detail regarding the vector unit **1410** and the vector registers **1414**. Specifically, the vector unit **1410** is a 16-wide vector processing unit (VPU) (see the 16-wide ALU **1428**), which executes one or more of integer, single-precision float, and double-precision float instructions. The VPU supports swizzling the register inputs with swizzle unit **1420**, numeric conversion with numeric convert units **1422A-B**, and replication with replication unit **1424** on the memory input. Write mask registers **1426** allow predicating resulting vector writes.

(228) FIG. **15** is a block diagram of a processor **1500** that may have more than one core, may have an integrated memory controller, and may have integrated graphics according to examples of the disclosure. The solid lined boxes in FIG. **15** illustrate a processor **1500** with a single core **1502A**, a system agent **1510**, a set of one or more bus controller units **1516**, while the optional addition of the dashed lined boxes illustrates an alternative processor **1500** with multiple cores **1502A-N**, a set of one or more integrated memory controller unit(s) **1514** in the system agent unit **1510**, and special purpose logic **1508**.

(229) Thus, different implementations of the processor **1500** may include: 1) a CPU with the special purpose logic **1508** being integrated graphics and/or scientific (throughput) logic (which may include one or more cores), and the cores **1502A-N** being one or more general purpose cores (e.g., general purpose in-order cores, general purpose out-of-order cores, a combination of the two); 2) a coprocessor with the cores **1502A-N** being a large number of special purpose cores intended primarily for graphics and/or scientific (throughput); and 3) a coprocessor with the cores **1502A-N** being a large number of general purpose in-order cores. Thus, the processor **1500** may be a general-purpose processor, coprocessor, or special-purpose processor, such as, for example, a network or communication processor, compression engine, graphics processor, GPGPU (general purpose graphics processing unit), a high-throughput many integrated core (MIC) coprocessor (including 30 or more cores), embedded processor, or the like. The processor may be implemented on one or more chips. The processor **1500** may be a part of and/or may be implemented on one or more substrates using any of a number of process technologies, such as, for example, BiCMOS, CMOS, or NMOS.

(230) The memory hierarchy includes one or more levels of cache within the cores, a set or one or more shared cache units **1506**, and external memory (not shown) coupled to the set of integrated memory controller units **1514**. The set of shared cache units **1506** may include one or more mid-level caches, such as level 2 (L2), level 3 (L3), level 4 (L4), or other levels of cache, a last level cache (LLC), and/or combinations thereof. While in one example a ring based interconnect unit

1512 interconnects the integrated graphics logic **1508**, the set of shared cache units **1506**, and the system agent unit **1510**/integrated memory controller unit(s) **1514**, alternative examples may use any number of well-known techniques for interconnecting such units. In one example, coherency is maintained between one or more cache units **1506** and cores **1502-A-N**.

(231) In some examples, one or more of the cores **1502A-N** are capable of multi-threading. The system agent **1510** includes those components coordinating and operating cores **1502A-N**. The system agent unit **1510** may include for example a power control unit (PCU) and a display unit. The PCU may be or include logic and components needed for regulating the power state of the cores **1502A-N** and the integrated graphics logic **1508**. The display unit is for driving one or more externally connected displays.

(232) The cores **1502A-N** may be homogenous or heterogeneous in terms of architecture instruction set; that is, two or more of the cores **1502A-N** may be capable of execution the same instruction set, while others may be capable of executing only a subset of that instruction set or a different instruction set.

(233) Exemplary Computer Architectures

(234) FIGS. **16-19** are block diagrams of exemplary computer architectures. Other system designs and configurations known in the arts for laptops, desktops, handheld PCs, personal digital assistants, engineering workstations, servers, network devices, network hubs, switches, embedded processors, digital signal processors (DSPs), graphics devices, video game devices, set-top boxes, micro controllers, cell phones, portable media players, hand held devices, and various other electronic devices, are also suitable. In general, a huge variety of systems or electronic devices capable of incorporating a processor and/or other execution logic as disclosed herein are generally suitable.

(235) Referring now to FIG. **16**, shown is a block diagram of a system **1600** in accordance with one example of the present disclosure. The system **1600** may include one or more processors **1610**, **1615**, which are coupled to a controller hub **1620**. In one example the controller hub **1620** includes a graphics memory controller hub (GMCH) **1690** and an Input/Output Hub (IOH) **1650** (which may be on separate chips); the GMCH **1690** includes memory and graphics controllers to which are coupled memory **1640** and a coprocessor **1645**; the IOH **1650** is couples input/output (I/O) devices **1660** to the GMCH **1690**. Alternatively, one or both of the memory and graphics controllers are integrated within the processor (as described herein), the memory **1640** and the coprocessor **1645** are coupled directly to the processor **1610**, and the controller hub **1620** in a single chip with the IOH **1650**. Memory **1640** may include code **1640A**, for example, to store code that when executed causes a processor to perform any method of this disclosure.

(236) The optional nature of additional processors **1615** is denoted in FIG. **16** with broken lines. Each processor **1610**, **1615** may include one or more of the processing cores described herein and may be some version of the processor **1500**.

(237) The memory **1640** may be, for example, dynamic random access memory (DRAM), phase change memory (PCM), or a combination of the two. For at least one example, the controller hub **1620** communicates with the processor(s) **1610**, **1615** via a multi-drop bus, such as a frontside bus (FSB), point-to-point interface such as Quickpath Interconnect (QPI), or similar connection **1695**.

(238) In one example, the coprocessor **1645** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like. In one example, controller hub **1620** may include an integrated graphics accelerator.

(239) There can be a variety of differences between the physical resources **1610**, **1615** in terms of a spectrum of metrics of merit including architectural, microarchitectural, thermal, power consumption characteristics, and the like.

(240) In one example, the processor **1610** executes instructions that control data processing operations of a general type. Embedded within the instructions may be coprocessor instructions.

The processor **1610** recognizes these coprocessor instructions as being of a type that should be executed by the attached coprocessor **1645**. Accordingly, the processor **1610** issues these coprocessor instructions (or control signals representing coprocessor instructions) on a coprocessor bus or other interconnect, to coprocessor **1645**. Coprocessor(s) **1645** accept and execute the received coprocessor instructions.

(241) Referring now to FIG. **17**, shown is a block diagram of a first more specific exemplary system **1700** in accordance with an example of the present disclosure. As shown in FIG. **17**, multiprocessor system **1700** is a point-to-point interconnect system, and includes a first processor **1770** and a second processor **1780** coupled via a point-to-point interconnect **1750**. Each of processors **1770** and **1780** may be some version of the processor **1500**. In one example of the disclosure, processors **1770** and **1780** are respectively processors **1610** and **1615**, while coprocessor **1738** is coprocessor **1645**. In another example, processors **1770** and **1780** are respectively processor **1610** coprocessor **1645**.

(242) Processors **1770** and **1780** are shown including integrated memory controller (IMC) units **1772** and **1782**, respectively. Processor **1770** also includes as part of its bus controller units point-to-point (P-P) interfaces **1776** and **1778**; similarly, second processor **1780** includes P-P interfaces **1786** and **1788**. Processors **1770**, **1780** may exchange information via a point-to-point (P-P) interface **1750** using P-P interface circuits **1778**, **1788**. As shown in FIG. **17**, IMCs **1772** and **1782** couple the processors to respective memories, namely a memory **1732** and a memory **1734**, which may be portions of main memory locally attached to the respective processors.

(243) Processors **1770**, **1780** may each exchange information with a chipset **1790** via individual P-P interfaces **1752**, **1754** using point to point interface circuits **1776**, **1794**, **1786**, **1798**. Chipset **1790** may optionally exchange information with the coprocessor **1738** via a high-performance interface **1739**. In one example, the coprocessor **1738** is a special-purpose processor, such as, for example, a high-throughput MIC processor, a network or communication processor, compression engine, graphics processor, GPGPU, embedded processor, or the like.

(244) A shared cache (not shown) may be included in either processor or outside of both processors, yet connected with the processors via P-P interconnect, such that either or both processors' local cache information may be stored in the shared cache if a processor is placed into a low power mode.

(245) Chipset **1790** may be coupled to a first bus **1716** via an interface **1796**. In one example, first bus **1716** may be a Peripheral Component Interconnect (PCI) bus, or a bus such as a PCI Express bus or another third generation I/O interconnect bus, although the scope of the present disclosure is not so limited.

(246) As shown in FIG. **17**, various I/O devices **1714** may be coupled to first bus **1716**, along with a bus bridge **1718** which couples first bus **1716** to a second bus **1720**. In one example, one or more additional processor(s) **1715**, such as coprocessors, high-throughput MIC processors, GPGPU's, accelerators (such as, e.g., graphics accelerators or digital signal processing (DSP) units), field programmable gate arrays, or any other processor, are coupled to first bus **1716**. In one example, second bus **1720** may be a low pin count (LPC) bus. Various devices may be coupled to a second bus **1720** including, for example, a keyboard and/or mouse **1722**, communication devices **1727** and a storage unit **1728** such as a disk drive or other mass storage device which may include instructions/code and data **1730**, in one example. Further, an audio I/O **1724** may be coupled to the second bus **1720**. Note that other architectures are possible. For example, instead of the point-to-point architecture of FIG. **17**, a system may implement a multi-drop bus or other such architecture.

(247) Referring now to FIG. **18**, shown is a block diagram of a second more specific exemplary system **1800** in accordance with an example of the present disclosure. Like elements in FIGS. **17** and **18** bear like reference numerals, and certain aspects of FIG. **17** have been omitted from FIG. **18** in order to avoid obscuring other aspects of FIG. **18**.

(248) FIG. **18** illustrates that the processors **1770**, **1780** may include integrated memory and I/O

control logic (“CL”) **1772** and **1782**, respectively. Thus, the CL **1772**, **1782** include integrated memory controller units and include I/O control logic. FIG. **18** illustrates that not only are the memories **1732**, **1734** coupled to the CL **1772**, **1782**, but also that I/O devices **1814** are also coupled to the control logic **1772**, **1782**. Legacy I/O devices **1815** are coupled to the chipset **1790**. (249) Referring now to FIG. **19**, shown is a block diagram of a SoC **1900** in accordance with an example of the present disclosure. Similar elements in FIG. **15** bear like reference numerals. Also, dashed lined boxes are optional features on more advanced SoCs. In FIG. **19**, an interconnect unit(s) **1902** is coupled to: an application processor **1910** which includes a set of one or more cores **1502A-N** and shared cache unit(s) **1506**; a system agent unit **1510**; a bus controller unit(s) **1516**; an integrated memory controller unit(s) **1514**; a set or one or more coprocessors **1920** which may include integrated graphics logic, an image processor, an audio processor, and a video processor; an static random access memory (SRAM) unit **1930**; a direct memory access (DMA) unit **1932**; and a display unit **1940** for coupling to one or more external displays. In one example, the coprocessor(s) **1920** include a special-purpose processor, such as, for example, a network or communication processor, compression engine, GPGPU, a high-throughput MIC processor, embedded processor, or the like.

(250) Examples (e.g., of the mechanisms) disclosed herein may be implemented in hardware, software, firmware, or a combination of such implementation approaches. Examples of the disclosure may be implemented as computer programs or program code executing on programmable systems comprising at least one processor, a storage system (including volatile and non-volatile memory and/or storage elements), at least one input device, and at least one output device.

(251) Program code, such as code **1730** illustrated in FIG. **17**, may be applied to input instructions to perform the functions described herein and generate output information. The output information may be applied to one or more output devices, in known fashion. For purposes of this application, a processing system includes any system that has a processor, such as, for example; a digital signal processor (DSP), a microcontroller, an application specific integrated circuit (ASIC), or a microprocessor.

(252) The program code may be implemented in a high level procedural or object oriented programming language to communicate with a processing system. The program code may also be implemented in assembly or machine language, if desired. In fact, the mechanisms described herein are not limited in scope to any particular programming language. In any case, the language may be a compiled or interpreted language.

(253) One or more aspects of at least one example may be implemented by representative instructions stored on a machine-readable medium which represents various logic within the processor, which when read by a machine causes the machine to fabricate logic to perform the techniques described herein. Such representations, known as “IP cores” may be stored on a tangible, machine readable medium and supplied to various customers or manufacturing facilities to load into the fabrication machines that actually make the logic or processor.

(254) Such machine-readable storage media may include, without limitation, non-transitory, tangible arrangements of articles manufactured or formed by a machine or device, including storage media such as hard disks, any other type of disk including floppy disks, optical disks, compact disk read-only memories (CD-ROMs), compact disk rewritables (CD-RWs), and magneto-optical disks, semiconductor devices such as read-only memories (ROMs), random access memories (RAMs) such as dynamic random access memories (DRAMs), static random access memories (SRAMs), erasable programmable read-only memories (EPROMs), flash memories, electrically erasable programmable read-only memories (EEPROMs), phase change memory (PCM), magnetic or optical cards, or any other type of media suitable for storing electronic instructions.

(255) Accordingly, examples of the disclosure also include non-transitory, tangible machine-

readable media containing instructions or containing design data, such as Hardware Description Language (HDL), which defines structures, circuits, apparatuses, processors and/or system features described herein. Such examples may also be referred to as program products.

(256) Emulation (Including Binary Translation, Code Morphing, Etc.)

(257) In some cases, an instruction converter may be used to convert an instruction from a source instruction set to a target instruction set. For example, the instruction converter may translate (e.g., using static binary translation, dynamic binary translation including dynamic compilation), morph, emulate, or otherwise convert an instruction to one or more other instructions to be processed by the core. The instruction converter may be implemented in software, hardware, firmware, or a combination thereof. The instruction converter may be on processor, off processor, or part on and part off processor.

(258) FIG. 20 is a block diagram contrasting the use of a software instruction converter to convert binary instructions in a source instruction set to binary instructions in a target instruction set according to examples of the disclosure. In the illustrated example, the instruction converter is a software instruction converter, although alternatively the instruction converter may be implemented in software, firmware, hardware, or various combinations thereof. FIG. 20 shows a program in a high level language **2002** may be compiled using an x86 compiler **2004** to generate x86 binary code **2006** that may be natively executed by a processor with at least one x86 instruction set core **2016**. The processor with at least one x86 instruction set core **2016** represents any processor that can perform substantially the same functions as an Intel® processor with at least one x86 instruction set core by compatibly executing or otherwise processing (1) a substantial portion of the instruction set of the Intel® x86 instruction set core or (2) object code versions of applications or other software targeted to run on an Intel® processor with at least one x86 instruction set core, in order to achieve substantially the same result as an Intel® processor with at least one x86 instruction set core. The x86 compiler **2004** represents a compiler that is operable to generate x86 binary code **2006** (e.g., object code) that can, with or without additional linkage processing, be executed on the processor with at least one x86 instruction set core **2016**. Similarly, FIG. 20 shows the program in the high level language **2002** may be compiled using an alternative instruction set compiler **2008** to generate alternative instruction set binary code **2010** that may be natively executed by a processor without at least one x86 instruction set core **2014** (e.g., a processor with cores that execute the MIPS instruction set of MIPS Technologies of Sunnyvale, CA and/or that execute the ARM instruction set of ARM Holdings of Sunnyvale, CA). The instruction converter **2012** is used to convert the x86 binary code **2006** into code that may be natively executed by the processor without an x86 instruction set core **2014**. This converted code is not likely to be the same as the alternative instruction set binary code **2010** because an instruction converter capable of this is difficult to make; however, the converted code will accomplish the general operation and be made up of instructions from the alternative instruction set. Thus, the instruction converter **2012** represents software, firmware, hardware, or a combination thereof that, through emulation, simulation or any other process, allows a processor or other electronic device that does not have an x86 instruction set processor or core to execute the x86 binary code **2006**.

Claims

1. An apparatus comprising: an execution circuit to execute an instruction that generates a memory access request for an element in memory via a first capability; a capability management circuit to check the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; a cache; and a prefetch circuit to: prefetch an additional element of the first object from the memory based on the first capability checked by the capability management circuit, determine if the additional

- element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access, and prefetch the second element from the memory to the cache based on the additional element being the second capability.
2. The apparatus of claim 1, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory to the cache based on the first capability.
 3. The apparatus of claim 1, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability.
 4. The apparatus of claim 3, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability.
 5. The apparatus of claim 3, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits are to prefetch the second element from the memory to the cache in parallel with the prefetch of the one or more further elements of the second object from the memory to the cache.
 6. The apparatus of claim 1, wherein the prefetch circuit is to: prefetch a second additional element of the first object from the memory based on the first capability, determine if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third object to which the third capability authorizes access, and prefetch the third element from the memory to the cache based on the second additional element being the third capability.
 7. The apparatus of claim 6, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the third object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the third capability.
 8. The apparatus of claim 1, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability.
 9. A method comprising: executing, by an execution circuit of a processor, an instruction that generates a memory access request for an element in memory via a first capability; checking, by a capability management circuit of the processor, the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; prefetching, by a prefetch circuit of the processor, an additional element of the first object from the memory based on the first capability checked by the capability management circuit; determining, by the processor, if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access; and prefetching, by the prefetch circuit, the second element from the memory to a cache of the processor based on the additional element being the second capability.
 10. The method of claim 9, wherein the prefetching of the additional element comprises prefetching the additional element of the first object from the memory to the cache based on the first capability.
 11. The method of claim 9, further comprising prefetching, by the prefetch circuit, one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability.
 12. The method of claim 11, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability.

13. The method of claim 11, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits perform the prefetching of the second element from the memory to the cache in parallel with the prefetching of the one or more further elements of the second object from the memory to the cache.

14. The method of claim 9, further comprising: prefetching, by the prefetch circuit, a second additional element of the first object from the memory based on the first capability; determining, by the processor, if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third object to which the third capability authorizes access; and prefetching, by the prefetch circuit, the third element from the memory to the cache based on the second additional element being the third capability.

15. The method of claim 14, further comprising prefetching, by the prefetch circuit, one or more further elements of the third object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the third capability.

16. The method of claim 9, wherein the prefetching of the additional element of the first object from the memory is based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability.

17. A system comprising: a memory; a processor comprising: an execution circuit to execute an instruction that generates a memory access request for an element in the memory via a first capability, and a cache separate from the memory; a capability management circuit to check the first capability for the memory access request, the first capability comprising an address field of the element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a first object to which the first capability authorizes access; and a prefetch circuit to: prefetch an additional element of the first object from the memory based on the first capability checked by the capability management circuit, determine if the additional element is a second capability comprising an address field of a second element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a second object to which the second capability authorizes access, and prefetch the second element from the memory to the cache based on the additional element being the second capability.

18. The system of claim 17, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory to the cache based on the first capability.

19. The system of claim 17, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the second object from the memory to the cache based on the one or more further elements being within the lower bound and the upper bound indicated by the second capability.

20. The system of claim 19, wherein the one or more further elements of the second object are a plurality of further elements of the second object within the lower bound and the upper bound indicated by the second capability.

21. The system of claim 19, wherein the prefetch circuit comprises a plurality of prefetch circuits, and the plurality of prefetch circuits are to prefetch the second element from the memory to the cache in parallel with the prefetch of the one or more further elements of the second object from the memory to the cache.

22. The system of claim 17, wherein the prefetch circuit is to: prefetch a second additional element of the first object from the memory based on the first capability, determine if the second additional element is a third capability comprising an address field of a third element in the memory, a validity field, and a bounds field that is to indicate a lower bound and an upper bound of a third object to which the third capability authorizes access, and prefetch the third element from the memory to the cache based on the second additional element being the third capability.

23. The system of claim 22, wherein the prefetch circuit is to cause a prefetch of one or more further elements of the third object from the memory to the cache based on the one or more further

elements being within the lower bound and the upper bound indicated by the third capability.

24. The system of claim 17, wherein the prefetch circuit is to prefetch the additional element of the first object from the memory based on the additional element being within the lower bound and the upper bound of the first object indicated by the first capability.
