

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417105
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Said; Farouk

Extensible platform for orchestration of data with built-in scalability and clustering

Abstract

In a computer system, an orchestration platform includes extensible components that interact with external systems and technology. The platform scales by way of a plurality of application servers using a clustering architecture.

Inventors:	Said; Farouk (Dubai, AE)
Applicant:	Keross FZ-LLC (Dubai, AE)
Family ID:	1000008819129
Assignee:	Keross FZ-LLC (Dubai, AE)
Appl. No.:	17/444763
Filed:	August 10, 2021

Prior Publication Data

Document Identifier	Publication Date
US 20230052148 A1	Feb. 16, 2023

Publication Classification

Int. Cl.: G06F9/44 (20180101); G06F9/445 (20180101); G06F9/54 (20060101)

U.S. Cl.:

CPC G06F9/44505 (20130101); G06F9/547 (20130101);

Field of Classification Search

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7490073	12/2008	Qureshi et al.	N/A	N/A
8539567	12/2012	Logue et al.	N/A	N/A
9392077	12/2015	Borzycki	N/A	G06F 21/6218
10120734	12/2017	Doraiswamy et al.	N/A	N/A
10181988	12/2018	Farah	N/A	N/A
11102251	12/2020	Pettit et al.	N/A	N/A
11669373	12/2022	Mittermeier	N/A	N/A
2006/0248189	12/2005	Jerrard-Dunne	709/225	H04L 63/168
2007/0192464	12/2006	Tullberg et al.	N/A	N/A
2010/0106678	12/2009	Pietrek et al.	N/A	N/A
2010/0333105	12/2009	Horvitz et al.	N/A	N/A
2012/0243416	12/2011	Hussain	726/11	H04L 12/40182
2012/0263440	12/2011	Malin	386/299	H04N 21/6581
2015/0378763	12/2014	Hassine et al.	N/A	N/A
2016/0162478	12/2015	Blassin	706/12	G06Q 10/063112
2016/0255420	12/2015	McClelland et al.	N/A	N/A
2017/0006113	12/2016	Singhal	N/A	N/A
2017/0034038	12/2016	Downey et al.	N/A	N/A
2017/0187573	12/2016	Haswell	N/A	N/A
2017/0364600	12/2016	Sadauskas, Jr.	N/A	G06F 16/31
2018/0295134	12/2017	Gupta et al.	N/A	N/A
2018/0322170	12/2017	Alberton et al.	N/A	N/A
2019/0215688	12/2018	Zavesky et al.	N/A	N/A
2019/0280949	12/2018	Wang	N/A	H04L 41/22
2020/0092258	12/2019	Stokes	N/A	N/A
2021/0090080	12/2020	Kang et al.	N/A	N/A
2021/0294596	12/2020	Ashirvad et al.	N/A	N/A
2021/0400085	12/2020	Stokes	N/A	N/A
2023/0050212	12/2022	Said	N/A	N/A

OTHER PUBLICATIONS

Non-final Office Action from U.S. Appl. No. 17/444,762, dated Aug. 6, 2023, 8 pp. cited by applicant

USPTO, Non-Final Rejection dated Feb. 1, 2024, in U.S. Appl. No. 17/444,761, 177704_17444761_02-01-2024_CTNE.pdf, Feb. 1, 2024. cited by applicant

USPTO, Final Rejection dated Jul. 25, 2024, in U.S. Appl. No. 17/444,761, 177704_17444761_07-25-2024_CTFR.pdf, Jul. 25, 2024. cited by applicant

USPTO, Advisory Action dated Oct. 16, 2024, in U.S. Appl. No. 17/444,761, 177704_17444761_10-16-2024_CTAV.pdf, Oct. 16, 2024. cited by applicant

USPTO, Non-Final Rejection dated Aug. 7, 2023, in U.S. Appl. No. 17/444,762, 177704_17444762_08-07-2023_CTNE.pdf, Aug. 7, 2023. cited by applicant

Primary Examiner: Pandey; Keshab R

Attorney, Agent or Firm: ESPE Legal Consultancy FZ-LLC

Background/Summary

TECHNICAL FIELD

(1) The invention pertains to computer systems, including orchestration of data reception, processing, storage, and retrieval.

BACKGROUND

(2) Computer systems increasingly manage large amounts of data on behalf of system users. In a typical environment, a large number of applications require access to data stored in a variety of databases with different structures and use cases. Over time, this diversity can create compatibility problems. For example, it may not be possible to make changes in one layer of a technology stack without having to make changes to the entire stack. Significant time and expense invested in a particular configuration can be lost when upgrading system components or migrating to new solutions.

(3) Known orchestration solutions attempt to solve some aspects of these problems by virtualizing the interface between applications and data. The added orchestration layer between applications and data makes it easier to manage changes in system applications and databases. But interoperability with applications that are outside the system remains an issue. There is a need for computer systems that integrate internal applications and data with outside systems and technology.

SUMMARY

(4) An orchestration platform includes an application stack and a database stack. The platform is extended by components that integrate and interact with technologies outside the core platform. The platform extension deploys a surrogate component or probe that acts as a bridge between the core platform and the extension technology. Multiple instances of the probe can be deployed on the platform extension in accordance with technical requirements.

(5) The probe is a standalone component that can be installed in Unix/Windows based computers, server or desktop with a unique key that is generated in the admin portal of the platform.

(6) In an embodiment, the probe executes Javascript, Jython or Scala scripts. The scripts can use the embedded libraries and application program interfaces (APIs) that the probe exposes. It can also download additional libraries (JAR files) on demand, to incorporate functionality not available off-the-shelf.

(7) In a further embodiment, the probe establishes connection with the platform and polls for “instructions” to be executed in the probe on schedule. The instructions are posted in the platform, tagged for a probe. These instructions contain “Remote Execution Service” definitions, i.e., scripts that use probe libraries and APIs to connect to any external systems or technology and send collected data to the platform.

(8) In an embodiment, the platform includes up to four application servers that can be horizontally scaled. One of these servers acts as a portal server. This server serves static content, such as html content, and bundled libraries, such as JavaScript libraries. A second server acts as a web server. The web server serves dynamic data as a REST request and response cycle. A third server acts as a job server. The job server executes background jobs as part of process models. A fourth server is a collector server that collects data from probes. Collector server nodes subscribe to various probe

data receiver queues, and the queues serve data in a round-robin strategy to ensure that only one collector server node processes the same data at the same time. The platform further provides semantics for distributed locking and distributed caching to manage clustered data processing environments as needed.

Description

SUMMARY OF FIGURES

- (1) FIG. 1 shows details of an orchestration platform ecosystem.
- (2) FIG. 2 shows interaction between the platform ecosystem and remote network premises with platform probes.
- (3) FIG. 3 shows details of interactions between an access control layer and a probe and user devices.
- (4) FIG. 4 shows details of interactions between the access control layer and the rest of the core platform ecosystem.
- (5) FIG. 5 shows logical-deployment details of interactions between probes associated with a platform extension and the core platform.
- (6) FIG. 6 shows physical-deployment details of interactions between an extension server and two platform servers.
- (7) FIG. 7 shows details of platform-extending probes in various configurations.
- (8) FIG. 8 shows details of application security for an orchestration platform.
- (9) FIG. 9 shows details of data and storage security for an orchestration platform.
- (10) FIG. 10 shows details of transport and network security for an orchestration platform.

DETAILED DESCRIPTION

- (11) An orchestration platform comprises an ecosystem that includes an application server stack and a database server stack. Additional platform extension architecture includes platform probes. Multiple instances of these probes can be installed at remote network locations. The platform probes may be controlled from an administration console of the platform.
- (12) In one embodiment, probes come with specialized libraries. The focus of these libraries follows particular use cases such as robotics, data collection from industry standard databases, data collection and management of IP-enabled devices, and remote program and script execution.
- (13) Extension libraries can also be injected into the probe after installation. These libraries enhance or upgrade existing probe capabilities to incorporate various technologies.
- (14) An HTTPD server acts as the entry point to the platform. This HTTPD server may also act as a request forwarder and load balancer to the platform application server stack. In the description that follows, HTTPD refers to Apache HyperText Transfer Protocol or Apache HTTP Server. Alternatively, nginx or another solution with similar functionality may be used.
- (15) FIG. 1 shows an orchestration-platform ecosystem **100**. Platform ecosystem **102** includes an application server stack **104**. Application server stack **104** includes job server **106**, and Message Queue (MQ or IBM MQ) server **108**. REST server **110** includes web servers **1** and **2** (**112**, **114**). HTML Content Server **116** is also part of application server stack **104**. An HTTPD server **118** includes load balancer **120** and request forwarder **122**. In communication with application server stack **104** are databases **124**, **126**, **128**, **130**, **132**, and **134**. Databases **124-134** are selected according to use cases and may include Cassandra, MongoDB, MySQL, MariaDB, ElasticSearch, and Redis.
- (16) FIG. 2 shows a detailed view **200** of the interaction between platform ecosystem **102** and an extension of the platform to a remote network. Compressed data in JSON or XML passes by way of connections **202** between platform ecosystem **102** and user devices **206**. Similarly, compressed data in JSON or XML passes between platform **102** and remote network premises **208** and **210**.

Remote network premises **208**, **210** each include a platform probe (**212**, **214**). Details of probes **212**, **214** will be described below.

(17) FIG. 3 shows details **300** of the interaction between access control layer **302**, which comprises REST facade **304** and WebSocket facade **306**, and user devices **206**. Communication between user devices **206** and access control layer **302** takes place by way of HTTP **308** and WebSocket **310**. Communication with access control layer **302** includes getting instructions **312** and sending data **314** between access control layer **302** and probe **316**.

(18) FIG. 4 shows detail **400** showing the interaction between access control layer **302** and the rest of the core platform ecosystem. Access control layer **302** communicates with administrative service **402**, application modeling service **404**, application runtime service **406** and data collection service **410**. Application modelling service **404** uses cache **412** for models, forms, and scripts. Cache **412** is used by application runtime service **406** which outputs to process engine **414** and MQ **416**, as well as to instruction service **408**. Job server **418** receives input from MQ **416** and communicates with instruction service **408**, which in turn interacts with access control layer **302**. Data collection service **410** communicates with MQ **416** and data store **420**. A command database (CMDB) **422** is accessible to administrative service **402**, application modelling service **404**, process engine **414**, and job server **418**.

(19) FIG. 5 shows detail **500** of communication between core platform **502** and platform extension **504**. Application server stack **104** communicates with probes **506** located outside the platform core. The database stack supporting the core platform is supplied by databases **124**, **126**, **128**, **130**, **132**, and **134**. The number and choice of databases varies depending on particular use cases. Examples of possible choices are Cassandra, MongoDB, MySQL, MariaDB, Elasticsearch, and Redis.

(20) Probes **506** include extension technologies **508**, **510**, **512**, **514**, and **516**. Examples of such technologies include Autodesk, ArcGIS, Node.js, R, and Spark. The probe exposes an environment of libraries and APIs to interact with any external systems, via various techniques, e.g., an in-process client for an external system like a database or a proprietary system, a facilitator for executing an R script to an adjacent R execution environment, or a node.js JavaScript to be executed in an adjacent node.js etc.

(21) FIG. 6 shows detail **600** of Unix platform server **602**, Unix platform server **604** and extension server **606**. Incoming communications using port **80** arrive at Unix platform server **602**. In a typical embodiment, platform servers **602** and **604** run an open source Unix-based Linux distribution such as Ubuntu while platform server **606** may be proprietary, such as Microsoft Windows. The actual choice of operating system for each of servers **602**, **604**, and **606** may be changed according to particular use cases.

(22) Platform **602** includes a MySQL database **608**, Redis database **610**, an application stack **612**, a load balancer, and an HTTPD server. Platform **604** supports the database stack and includes databases **616**, **618**, **620**, **622**, and **624**. On extension platform server **606** reside probes **626**, **628**, **630**, **632**, and **634**.

(23) One aspect of a probe of the present invention is that it executes scripts, for example, using programming languages such as Javascript, Jython, or Scala. Other languages may also be used, depending on particular use cases. The scripts use the embedded libraries and APIs that the probe exposes. The probe also downloads additional libraries, such as JAR files, to add additional functionality. A JAR file is a Java archive file format based on the ZIP file format that is used for aggregating many files into one.

(24) Another aspect of the probe is that it establishes connection with the platform and polls for “instructions” to be executed in the probe on schedule. The instructions are posted in the platform, tagged for a probe. These instructions contain “Remote Execution Service” definitions, i.e., scripts that use probe libraries and APIs to connect to any external systems or technology and send collected data to the platform.

(25) The probe is a standalone software component that sits in remote premises, in the vicinity of

the target systems it is configured to connect to. It builds a conduit to the platform to execute specific connector instructions. Software connectors transfer control and data among system components. For example, MariaDB Connector/Node.js is a native Javascript driver used to connect applications developed on Node.js to MariaDB and MySQL databases. Connectors also provide services that are independent of the interacting components' functionalities. Examples of such services are persistence, invocation, messaging and transactions. These services are sometimes known as “facilities components” by middleware standards such as CORBA, DCOM and RMI.

(26) In an embodiment, the probe is a framework for connectors, rather than an in-built connector. Hence the probe exposes an environment to execute scripts to inter-operate with an external system. The probe can load additional client libraries on-demand to connect to proprietary technologies. The probe, being a standalone component in the customer premises, its environment is not opaque as the platform and can be boosted with other software components. In an embodiment, all connectors in the probe are scripts to facilitate specific handling of data. Data collected and curated by the probe is sent to the platform, where it can be further manipulated in a historical context. An appropriate datastore is chosen depending on the nature of the data. In an embodiment, Influxdb or Cassandra are chosen for time-series data, Redis for geodesic data, and MariaDB ColumnStore for huge amounts of structured relational data. As new datastores are developed, they may be chosen using the same or similar criteria.

(27) FIG. 7 shows detail **700** of platform **702**, an embodiment of platform server **602** described above. In this embodiment, one or more probes are configured to carry out specific tasks. These specific tasks could be duplicated by one or more probes or distributed among the probes in various combinations. Illustrative examples of tasks performed by the probes include platform or technology extension, client data integration, fetching data from third-party providers, remote monitoring and management, or receiving data on-demand.

(28) First probe **704** is configured for and includes platform extension or technology integration **706** such as Autodesk, ArcGIS, Node.js, R, and Spark.

(29) Second probe **708** communicates with integrated client data **710**. In an embodiment, this integrated client data **710** includes connected applications or databases via APIs. Alternatively, integrated client data **710** is a subscribed message queue that uses, for example, the Advanced Message Queuing Protocol (AMQP) or the Message Queuing Telemetry Transport (MQTT). In another embodiment, integrated client data **710** comprises web scraping or desktop applications with Robotic Process Automation (“RPA”). In this context, RPA generally refers to software robotics that automate business-process activities.

(30) Third probe **712** is linked to a third-party data provider **714**. In an embodiment, data retrieved from the third party is accessed by third probe **712** but not replicated in other parts of the platform.

(31) Fourth probe **716** is configured for remote monitoring or management of one or more resources **718**. Exemplary monitored or managed resources **718** include a router, a firewall, a hub or router, mobile devices, laptop computers, Internet Protocol telephones, and websites. A system or network comprising different combinations of these resources is monitored or managed by fourth probe **716**.

(32) Probe n **720** receives data on-demand from data source **722**. Data source **722** alternatively comprises one or more kinds of internet sockets. In an embodiment, the sockets comprise raw User Transmission Control Protocol (TCP)/User Datagram Protocol (UDP) sockets. Alternatively, source **722** comprises Simple Object Access Protocol (SOAP) or Representational State Transfer (REST) interfaces. Source **722** may also comprise custom HTTP servers.

(33) FIG. 8 shows detail **800** showing security aspects of platform **802**. Remote devices **804**, such as desktop computers, laptop computers, tablets, and mobile phones, communicate with platform **802** by way of secure hypertext transfer protocol (HTTPS) **806**. Authenticated users **808** are given access by way of role-based-access-control (RBAC) to one or more accounts. In an embodiment,

available accounts are a first account **812** with applications (**814**, **816**) and a second account **818** with applications (**820**, **822**) and additional accounts represented by account **824**. The accounts **812**, **818**, through **824** are accessible to authenticated users. These accounts in turn have access to platform services **828**.

(34) Applications **814**, **816**, **820**, and **822** may receive dedicated service from RBAC **810**. For example, in an embodiment application **814** has on-demand security. In another embodiment, application **820** has on-demand data classification.

(35) FIG. **9** shows detail **900** of security aspects of platform **902**. In an embodiment, platform **902** is the same platform as platform **802** in FIG. **8**. Remote devices **904**, such as desktop computers, laptop computers, tablets, and mobile phones, communicate with platform **902** by way of secure hypertext transfer protocol (HTTPS) **906** to access application server stack **908**. The application server stack in turn communicates with one or more databases **912**.

(36) Data passing from application server **908** is encrypted by process **910**. In an embodiment, process **910** uses SHA2 Encryption.

(37) FIG. **10** shows detail **1000** of security aspects of platform **1002**. In an embodiment, platform **1002** is the same platform as platform **802** in FIG. **8** or the same platform as platform **902** in FIG. **9**. Remote devices **1004**, such as desktop computers, laptop computers, tablets, and mobile phones, communicate with platform **1002** by way of secure hypertext transfer protocol (HTTPS) **906** to access a cross-platform web server **1008**, such as Apache HTTPD server. Load balancer **1010** and request forwarder **1012** handle communication between server **1008** and application server stack **1014**. In this configuration, transport and network security is provided by HTTPS.

(38) The platform has a cluster architecture for dividing user requests among platform resources, such that a single user request can be handled and delivered by multiple server nodes.

(39) In an embodiment, the platform includes up to four application servers. In a further embodiment, these application servers are configured to be horizontally scaled. One of these servers acts as a portal server. This server serves static content, such as html content, and bundled libraries, such as JavaScript libraries. In a further embodiment, additional portal server nodes are added, and reverse proxied in the Apache HTTPD server which fronts all requests. A reverse proxy is a configuration where a server is positioned in front of web servers and forwards client requests to those web servers.

(40) In an embodiment, the platform includes a web server. The web server serves dynamic data as a REST request and response cycle. In a further embodiment, additional web server nodes are added, and reverse proxied by the Apache HTTPD server.

(41) In an embodiment, the platform includes a job server. The job server executes background jobs as part of process models. In a further embodiment, additional nodes of job server are added to the application server stack. Job server nodes are idempotent, such that only one node executes a job at the same time, by acquiring a lock on the persistent job store.

(42) In an embodiment, the platform includes a collector server that collects data from probes. In a further embodiment, additional nodes of the collector server are added to the application server stack. Collector server nodes subscribe to various probe data receiver queues, and the queues serve data in a round-robin strategy. Round robin refers generally to rotating requests among web servers in the order the requests are received. This strategy ensures that only one collector server node processes the same data at the same time.

(43) The platform further provides semantics for distributed locking and distributed caching to manage clustered data processing environments as needed. Distributed locking is a technique that ensures that two processes cannot both access shared data at the same time. The locking protocol ensures that only one process is allowed to proceed once a lock is established. In distributed caching, user data is not stored in the individual web server's memory, but on other available resources. Cached data is accessible to an application's web servers or virtual machines. The cached

data remains accessible to every server that runs the application, even when the application scales by adding or removing servers, or when servers are replaced due to upgrades or faults.

Claims

1. A computer system platform for data orchestration comprising: a platform core with an application stack and a database stack; an extension server, remote from the platform core; a probe comprising a standalone software component, installed on the extension server and in communication with the platform core by way of one or more connector; wherein the extension server with the installed probe is located proximate to a target system and wherein the target system is external to the platform core and comprises a datastore and a library corresponding to a programming language; the probe is configured to execute scripts that request data from the datastore in the programming language of the target system; a plurality of application servers in communication with the platform core, including a collector server; and wherein the target system further comprises an application program interface (“API”) connected to the library and the probe is configured to expose the API to make the library available to the platform core on-demand and receive data from the datastore at the platform core via the collector server.
2. The system of claim 1 wherein the plurality of application servers further includes at least two of a portal server, web server, and a job server.
3. The system of claim 1 wherein the plurality of application servers includes a portal server, web server, job server, and the collector server.
4. The system of claim 1 wherein the plurality of application servers includes a portal server in reverse proxy configuration, a web server, job server, and the collector server.
5. The system of claim 1 wherein the plurality of application servers includes a portal server in a reverse proxy configuration.
6. The system of claim 1 wherein the plurality of application servers includes a web server in a reverse proxy configuration.
7. A method of orchestrating data in a computer system comprising a platform core with an application stack and a database stack, an extension server, and an external target system, the method comprising: installing a probe comprising a standalone software component on the extension server at a remote location; configuring the probe to connect with the platform core, wherein the probe is installed on the extension server is in communication with the platform core by way of one or more connectors; wherein the extension server with the installed probe is located proximate to the external target system and wherein the external target system is external to the platform core; wherein the external target system comprises a datastore and a library corresponding to a programming language; wherein the probe is configured to execute scripts that request data from the datastore in the programming language of the external target system; configuring a plurality of application servers in communication with the platform core, including a collector server; and wherein the target system comprises an application program interface (“API”) connected to the library and the probe is configured to expose the API and the library to the platform core on-demand and receive data from the datastore at the platform core via the collector server.
8. The method of claim 7 wherein the plurality of application servers further includes at least two of a portal server, web server, and a job server.
9. The method of claim 7 wherein the plurality of application servers includes a portal server, web server, job server, and the collector server.
10. The method of claim 7 wherein the plurality of application servers includes a portal server in reverse proxy configuration, a web server, job server, and the collector server.
11. The method of claim 7 wherein the plurality of application servers includes a portal server in a reverse proxy configuration.

12. The method of claim 7 wherein the plurality of application servers includes a web server in a reverse proxy configuration.
13. A computer system platform for data orchestration, the system comprising: an extension server, remote from a platform core; a probe comprising a standalone software component, installed on the extension server and in communication with the platform core by way of one or more connectors; wherein the extension server with the installed probe is located proximate to an external target system and wherein the external target system remote from the platform core and comprises a datastore and a library corresponding to a programming language; wherein the probe is configured to execute scripts that request data from the datastore in the programming language of the external target system; a plurality of application servers in communication with the probe, including a collector server; and wherein the external target system comprises an application program interface (“API”) connected to the library and the probe is configured to expose the API and the library to the platform core on-demand and receive data from the datastore at the platform core via the collector server.
14. The system of claim 13 wherein the plurality of application servers further includes at least two of a portal server, web server, and a job server.
15. The system of claim 13 wherein the plurality of application servers includes a portal server, web server, job server, and the collector server.
16. The system of claim 13 wherein the plurality of application servers includes a portal server in reverse proxy configuration, web server, job server, and the collector server.
17. The system of claim 13 wherein the plurality of application servers includes a portal server in a reverse proxy configuration.
18. The system of claim 13 wherein the plurality of application servers includes a web server in a reverse proxy configuration.
19. The system of claim 17 wherein the plurality of application servers further includes a web server in a reverse proxy configuration.
20. The system of claim 17 wherein collector server comprises nodes subscribed to probe data receiver queues, and the queues are configured to serve round-robin data to the platform core via the collector server.
-