

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417155
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Chakravarthi Konduru; Dileep et al.

Failure recovery in a fault tolerant architecture

Abstract

A system/process provides fault tolerance to an integrated component system which integrates core components of a transaction processing system into a single processing platform, i.e., a single server, enabling elimination of the network interconnects and associated latencies introduced thereby in favor of much faster interconnects, such as inter-process communication and shared memory communication messaging, where a failure of any one component necessitates failing over the entire system to a backup thereof.

Inventors: Chakravarthi Konduru; Dileep (Naperville, IL), Soni; Priteshkumar (Naperville, IL), Koziol; Arkadiusz (Glenview, IL), Coffey; Matthew (Chicago, IL), Reddy; Nagesh (Elm Grove, WI), Kamineni; Satya (Lockport, IL), Sirohi; Madhur (Aurora, IL)

Applicant: Chicago Mercantile Exchange Inc. (Chicago, IL)

Family ID: 1000008818631

Assignee: Chicago Mercantile Exchange Inc. (Chicago, IL)

Appl. No.: 18/399922

Filed: December 29, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20250217244 A1	Jul. 03, 2025

Publication Classification

Int. Cl.: G06F11/20 (20060101); G06F11/14 (20060101)

U.S. Cl.:

Field of Classification Search

CPC: G06F (11/2023); G06F (11/1443); G06F (11/1464)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
4228496	12/1979	Katzman et al.	N/A	N/A
4590554	12/1985	Glazer et al.	N/A	N/A
5022076	12/1990	Rosenow	705/64	G06Q 20/382
5363503	12/1993	Gleeson	N/A	N/A
5621885	12/1996	Del, Jr.	N/A	N/A
6324654	12/2000	Wahl et al.	N/A	N/A
6691257	12/2003	Suffin	714/43	G06F 11/2007
6954877	12/2004	Earl et al.	N/A	N/A
6978396	12/2004	Ruuth et al.	N/A	N/A
7434096	12/2007	Callaway et al.	N/A	N/A
7480827	12/2008	Callaway et al.	N/A	N/A
8191078	12/2011	Cullen	714/4.11	H04L 67/1095
8295813	12/2011	Aalto	455/412.1	H04L 67/1001
9691102	12/2016	Studnitzer et al.	N/A	N/A
2003/0009707	12/2002	Pedone et al.	N/A	N/A
2007/0180302	12/2006	Allen	714/E11.08	G06F 11/2071
2008/0126853	12/2007	Callaway	714/13	G06F 11/1695
2010/0017647	12/2009	Callaway	714/48	G06F 11/1641
2010/0100475	12/2009	Callaway	705/37	G06F 11/1641
2011/0041006	12/2010	Fowler	N/A	N/A
2016/0110263	12/2015	Tracht	714/6.23	G06F 11/1435
2017/0123929	12/2016	Helleren	N/A	G06F 9/5088
2018/0167492	12/2017	Bonig	N/A	G06Q 40/04
2019/0391741	12/2018	Barinov	N/A	G06F 11/2038
2020/0050523	12/2019	Pawlowski	N/A	G06F 13/4022

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
1728697	12/2005	CN	N/A
2439881	12/2011	EP	H04L 67/1002
WO-2017023244	12/2016	WO	G06F 11/07

OTHER PUBLICATIONS

W. Zhao, L. E. Moser and P. M. Melliar-Smith, “Unification of replication and transaction processing in three-tier architectures,” Proceedings 22nd International Conference on Distributed Computing Systems, Vienna, Austria, 2002, pp. 290-297 (Year: 2002). cited by examiner

G. Su and A. Iyengar, “A highly available transaction processing system with non-disruptive failure handling,” 2012 IEEE Network Operations and Management Symposium, Maui, HI, USA, 2012, pp. 409-416 (Year: 2012). cited by examiner

Kumari et al., “A survey of fault tolerance in cloud computing”, Journal of King Saud University—Computer and Information Sciences 33 (2021) 1159-1176, available online Oct. 1, 2018, 18 pages. cited by applicant

Extended European Search Report in European Patent Application No. 24218848.0, dated May 19, 2025, 14 pages. cited by applicant

Pedone et al: “Pronto: High availability for standard off-the-shelf databases”, Journal of Parallel and Distributed Computing, Elsevier, Amsterdam, NL, vol. 68, No. 2, Jan. 5, 2008 (Jan. 5, 2008), pp. 150-164, XP022410959, ISSN: 0743-7315, DOI: 10.1016/J.JPDC.2007.08.001. cited by applicant

Primary Examiner: Kim; Matthew M

Assistant Examiner: Chowdhury; Indranil

Attorney, Agent or Firm: Lempia Summerfield Katz LLC

Background/Summary

CROSS REFERENCE TO RELATED APPLICATIONS

(1) The present patent application is related to co-pending U.S. patent application Ser. No. 18/399,917, filed contemporaneously herewith, entitled “FAULT TOLERANT ARCHITECTURE,” the entire disclosure of which is incorporated by reference herein.

BACKGROUND

(2) Fault Tolerance is generally regarded as the ability to mask, or recover from, erroneous conditions in a system once an error has been detected. Fault tolerance is typically required for mission critical systems/applications. Mission critical typically refers to any indispensable operation that cannot tolerate intervention, compromise, or shutdown during the performance of its critical function, e.g., any computer process that cannot fail during normal business hours. Examples of mission critical environments include business-essential process control, finance, such as the electronic trading systems described herein, health, safety, and security. These environments typically monitor, store, support and communicate data that cannot be lost or corrupted without compromising their core function.

(3) With respect to financial applications and electronic trading in particular, in addition to increased capacity and lower latency, the global nature of electronic trading systems has further driven a need for fault tolerance to increase their availability and reliability. Consistent reliable

operation may be critical to ensuring market stability, reliability, and acceptance. Therefore, scheduled outages should be minimized and unscheduled outages should be eliminated.

(4) In particular, critical applications/systems such as electronic trading systems often feature failure/disaster recovery mechanisms which allows the most current pre-failure state of a primary instance of the application/system, e.g., the state of the database thereof, to be recovered, restored, replicated and/or restarted, etc. in the event of an otherwise uncorrectable or unrecoverable temporary or permanent failure of one or more components of the primary instance of the application/system or the entirety thereof, or one or more components of the infrastructure upon which the primary instance is implemented. The state of the system often refers to the state of the most recent committed transaction prior to the failure.

(5) These failure recovery mechanisms may take the form of a backup component provided for each of one or more components of the primary instance. One type of backup may provide a fully redundant copy/instance of each of the one or more components of the primary instance, including any databases used thereby, which receives copies of all inputs to the primary instance and processes those inputs upon receipt and stores the results just as the primary instance would but, for example, reserving its outputs rather than sending them on to the consumers of those outputs, etc., e.g., to avoid delivering duplicates of the outputs of the primary instance during normal operations. While potentially more expensive, with this type of backup system, when a component of the primary instance fails, the backup component can simply be switched over to be used in its place, with some caveats. This may provide for a faster and/or more reliable recovery after a failure occurs. Another type of backup may simply comprise a database or other storage which stores copies of the inputs to the primary instance along with the outputs generated thereby or otherwise a log of which inputs were successfully processed by the primary instance, referred to as a storage-based backup system, such that when the primary instance fails and a backup instance is started, the backup store need only be consulted to determine which inputs were received but not completely processed by the primary instance, i.e., due to the failure, and therefore need to be re-processed by the backup instance before it can assume normal processing of newly received inputs.

(6) When implementing a backup system/instance, the operator needs to consider where to locate, physically and/or logically, that backup system relative to the primary instance and the source of inputs thereto. For example, in the case of physical infrastructure issues or natural disasters which may affect the geographic location of the primary instance, one would not want the backup system to be physically located in the same geographic areas as then it may be vulnerable to the same issues, thereby reducing the effectiveness of the backup system. Accordingly, most operators locate their backup systems in different geographic regions from where the primary instance is located to minimize the chance that both systems will experience failures due to the same cause. It will be appreciated that many operators deploy multiple backup systems located in disparate geographic regions to further minimize the likelihood of the same event compromising both the primary instance and the availability of at least one backup system.

(7) However, locating the backup system in a geographic location different from where the primary instance is located creates a latency issue with regard to communication of the transactional inputs to both systems to assure that the backup system is synchronized with the primary system should the backup system need to take over, i.e., to minimize transaction reprocessing or loss. One solution may be to receive the transactional inputs from their source at a location that is equidistant, or otherwise subject to substantially equal/symmetric communications latencies, to both the primary instance and the backup system, where the transactions are then relayed to both systems from this equidistant location. However, for storage-based backup systems, checkpoints from the primary instance may still need to be periodically communicated between the primary instance and the backup system to minimize the extent to which inputs need to be reprocessed when a failure occurs. Furthermore, in latency sensitive applications, such as financial transaction processing, e.g., electronic trading, it may be necessary to locate the primary instance close to the source of

transactions sent thereto so as to minimize the operational latency of the application/system. This may necessarily mean the backup system has to be located further from that source of transactions. (8) Accordingly, it is desirable to provide for fault tolerance of a high performance electronic trading system which minimizes recovery time and performance degradation in fail over scenarios.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 depicts an exchange computer and network system, according to some embodiments.
- (2) FIG. 2 depicts a general computer system, according to some embodiments.
- (3) FIG. 3 depicts a block diagram of a component integrated and fault tolerant system according to one embodiment.
- (4) FIG. 4 depicts a transaction processing server according to one embodiment.
- (5) FIG. 5 depicts a block diagram of a component integrated system and fault tolerant system according to one embodiment.
- (6) FIG. 6 depicts a block diagram of the inbound message flow into the primary instance according to one embodiment.
- (7) FIG. 7 depicts a block diagram of the input message flow into the backup instance according to one embodiment.
- (8) FIG. 8 depicts a block diagram of the outbound message flow from the primary instance according to one embodiment.
- (9) FIG. 9 depicts a block diagram of the outbound message flow from the backup instance according to some embodiments.
- (10) FIGS. 10A and 10B depict a sequence diagram of a startup sequence of a primary instance according to one embodiment.
- (11) FIGS. 11A-11D depict a sequence diagram of a startup sequence of a backup instance according to one embodiment.
- (12) FIGS. 12A(1), 12A(2), 12A(3), and 12B depict a sequence diagram of a failover sequence of a backup instance in a fault tolerant system.
- (13) FIG. 13 depicts a flowchart of the operation of an integrated component system according to one embodiment.
- (14) FIG. 14 depicts a flowchart of the operation of a fault tolerant system according to one embodiment.
- (15) FIG. 15 depicts a flowchart of the operation of an integrated component and fault tolerant system with different state transitions.

DETAILED DESCRIPTION

(16) The disclosed embodiments relate to a system/process that provides fault tolerance to an integrated component system which integrates core components of a transaction processing system into a single effective processing platform, e.g., a single or multi-processor/core server or tightly coupled set of single and/or multi-processor/core servers, enabling elimination of the network interconnects and associated latencies introduced in typical distributed processing systems in favor of much faster interconnects, such as inter-process communication, direct interconnect and shared memory communication messaging, where a failure of any one component typically necessitates failing over the entire system to a backup thereof.

(17) As will be discussed, the integrated implementation of the disclosed transaction processing system implicates the ability to provide redundancy for any one component thereof as the use of component level backup instances, typically, as described above, located remote from the primary instance to assure the reliability/availability thereof, would require, in a failover situation, that the correctly operating components now communicate with the backup instance of the failed

component over a network or other interconnect which necessarily has a larger communications latency with the still functioning components than the failed component had, necessarily impacting performance. Furthermore, in an integrated implementation many types of faults which might impact one component are likely to impact all components, e.g., compromises to operating power or network connectivity.

(18) A financial instrument trading system, such as a futures exchange, referred to herein also as an “Exchange”, such as the Chicago Mercantile Exchange Inc. (CME), provides a contract market where financial instruments, for example futures, options on futures and spread contracts, are traded among market participants, e.g., traders, brokers, etc.

(19) Current financial instrument trading systems allow traders to submit orders and receive confirmations, market data, and other information electronically via a communications network. These “electronic” marketplaces, implemented by, and also referred to as, “electronic trading systems,” are an alternative trading forum to pit based trading systems whereby the traders, or their representatives, all physically stand in a designated location, i.e., a trading pit, and trade with each other via oral and visual/hand based communication.

(20) Typically, the Exchange provides for centralized “clearing” by which all trades are confirmed and matched, and open positions are settled each day until expired (such as in the case of an option), offset or delivered. Matching, which is a function typically performed by the Exchange, is a process, for a given order which specifies a desire to buy or sell a quantity of a particular instrument at a particular price, of seeking/identifying one or more wholly or partially, with respect to quantity, satisfying counter orders thereto, e.g. a sell counter to an order to buy, or vice versa, for the same instrument at the same, or sometimes better, price (but not necessarily the same quantity), which are then paired for execution to complete a trade between the respective market participants (via the Exchange) and at least partially satisfy the desired quantity of one or both of the order and/or the counter order, with any residual unsatisfied quantity left to await another suitable counter order, referred to as “resting.”

(21) In particular, electronic trading of financial instruments, such as futures contracts, is conducted by market participants sending trading orders, such as to buy or sell one or more futures contracts, in electronic form to the Exchange. These electronically submitted orders to buy, and sell are then matched, if possible, by the Exchange, i.e., by an Exchange's Transaction Processor (TP), also referred to as a match engine or matching engine, to execute a trade, with the results thereof being communicated to the market participants through electronic notifications/broadcasts, referred to as market data feeds. Outstanding (unmatched, wholly unsatisfied/unfilled, or partially satisfied/filled) orders are maintained in one or more data structures or databases referred to as “order books,” such orders being referred to as “resting,” and made visible, i.e., their availability for trading is advertised, to the market participants through the electronic notifications/broadcasts, i.e., market data feeds, as well. An order book is typically maintained for each product, e.g., instrument, traded on the electronic trading system and generally defines or otherwise represents the state of the electronic trading system and of the market for that product, i.e., the current prices at which the market participants are willing buy or sell that product. As such, as used herein, an order book for a product may also be referred to as a market for that product.

(22) A market data feed, referred to as market data or market feed, is a compressed or uncompressed real time (with respect to market events), or substantial approximation thereof, an electronic data/message stream provided via an electronic communications network, such as the Internet, by the Exchange directly, or via a third party intermediary. A market data feed may be comprised of individual electronic messages, each comprising one or more packets or datagrams, and may carry, for example, pricing or other information regarding orders placed, traded instruments and other market information, such as summary values and statistical values, or combinations thereof, and may be transmitted, e.g., multi-casted, to the market participants using standardized protocols, such as UDP over Ethernet. More than one market data feed, each, for

example, carrying different information, may be provided. The standard protocol that is typically utilized for the transmission of market data feeds is the Financial Information Exchange (FIX) protocol Adapted for Streaming (FAST), aka FIX/FAST, which is used by multiple exchanges to distribute their market data. Pricing information conveyed by the market data feed may include the prices, or changes thereto, of resting orders, prices at which particular orders were recently traded, or other information representative of the state of the market or changes therein. Separate, directed/private, messages may also be transmitted directly to market participants to confirm receipt of orders, cancellation of orders and otherwise provide acknowledgment or notification of matching and other events relevant, or otherwise privy, only to the particular market participant.

(23) The GLOBEX® electronic trading system, offered by CME implements an electronic trading system/marketplace for trading futures and options (option contracts), referred to as Exchange Traded Derivative (ETD) options, on futures wherein the underlying is a futures contract for a particular underlier. They are listed and traded by Strike price and Expiry (daily, weekly, monthly, quarterly). ETD options physically expire into, i.e., upon expiration the contract delivers, the closest expiring future contract (typically a highly liquid, if not the most liquid, future contract) for the particular underlier, e.g., in the case of an ETD FX option, it is the closest quarterly expiring future contract. Then, the futures contract is settled physically or via cash.

(24) In particular, GLOBEX® is an open access marketplace that allows participants to directly enter their own trades and participate in the trading process, including viewing the book of orders and real-time price data. GLOBEX® has a number of core components/applications/engines or components including a transaction receiver processor (TR), e.g., a market segment gateway, a transaction processor (TP), e.g., a matching engine, a result generator (RG), e.g., a market data generator, and a transaction logger (TL) which includes a database that stores records for transactions for reporting, audit and historical purposes, each of which, in prior implementations, may be deployed in a distributed fashion, e.g., on different inter-networked physical servers. As will be described, in the disclosed embodiments, these components may be integrated, e.g., into a single system/server or tightly coupled set thereof.

(25) The following sequence describes how, at least in part, information may be propagated in an electronic trading system such as GLOBEX®, through a series of electronic messages, and how orders may be processed: (1) An opportunity is created at the TP, such as by placing a recently received but unmatched order on the order book to rest indicating a trader willing to trade a particular quantity at a particular price, i.e., an opportunity for another trader to trade therewith; (2) The TP creates an update reflecting the opportunity and sends it to the RG via the network; (3) The RG generates market data message based on the update and transmits it, e.g., using multicasting, or multiple unicast streams, or other mode of broadcasting via the network, the generated market data message to all of the market participants to advertise the opportunity to trade, i.e., transmits a broadcast market data feed message; (4) The market participants evaluate the opportunity and each, upon completion of their evaluation, may or may not choose to respond with an order, i.e., a message containing their counter-order, responsive to the resting order, i.e., counter to the resting order for some or all of the available quantity, by transmitting the counter-order via the network to the TR; (5) The TR receives any counter orders generated by the market participants, ascribes an ordering to each received order relative to the other received orders, sends confirmation of receipt back directly to each submitting market participant, and forwards the received orders to the TP via the network; (6) The TP evaluates the received orders and matches, using for example, a price/time priority algorithm, the first suitable arriving order, based on the ascribed ordering, against the resting opportunity and a trade is executed. Depending upon the available quantity of the resting order and the quantity of the incoming order, one of the resting order or the incoming order may be fully satisfied with any remainder of either order being placed on the order book to rest, or cancelled, etc. (7) The TP creates an update and sends it to the RG via the network. (8) The RG generates a market data message based on the update and transmits the generated market data

message to all of the market participants via the network.

(26) At various points in the above process, data regarding the processing of incoming orders and the final and/or intermediate results thereof may be stored by the TL.

(27) As will be appreciated, since different components may be implemented in one or more separate servers interconnected via a network connection, the latency of the messages communicated therebetween may impact the performance of the system.

(28) As described above, Fault Tolerance is generally regarded as the ability to mask, or recover from, erroneous conditions in a system once an error has been detected. Fault tolerance is typically required for mission critical systems/applications. Mission critical typically refers to any indispensable operation that cannot tolerate intervention, compromise, or shutdown during the performance of its critical function, e.g., any computer process that cannot fail during normal business hours. Examples of mission critical environments include business-essential process control, finance, such as the electronic trading systems described above, health, safety, and security. These environments typically monitor, store, support and communicate data that cannot be lost or corrupted without compromising their core function.

(29) In addition to increased capacity and lower latency, as already indicated above, the global nature of electronic trading systems has further driven a need for fault tolerance to increase their availability and reliability. Consistent reliable operation may be critical to ensuring market stability, reliability, and acceptance. Therefore, scheduled outages should be minimized and unscheduled outages should be eliminated.

(30) In particular, as described above, critical applications/systems such as electronic trading systems often feature failure/disaster recovery mechanisms which allows the state of a primary instance of the application/system, e.g., the state of the database thereof, to be recovered, restored, replicated and/or restarted, etc. in the event of an otherwise uncorrectable or unrecoverable temporary or permanent failure of one or more components of the primary instance of the application/system or the entirety thereof, or one or more components of the infrastructure upon which the primary instance is implemented. The state of the system often refers to the state of the most recent committed transaction prior to the failure.

(31) As noted above, these failure recovery mechanisms may take the form of a backup component for each component of the primary system. One type of backup may provide a fully redundant copy/instance of the primary component, including the primary component's database, which receives copies of all inputs to the primary instance and processes those inputs upon receipt and stores the results just as the primary system would but, for example, reserving its outputs rather than sending them on to the consumers of those outputs, etc., e.g., to avoid delivering duplicates the outputs of the primary instance during normal operations. While potentially more expensive, with this type of backup system, when the primary component fails, the backup component can simply be switched over to be used in its place, with some caveats. This may provide for a faster and/or more reliable recovery after a failure occurs. Another type of backup may simply comprise a database or other storage which stores copies of the inputs to the primary instance along with the outputs generated thereby or otherwise a log of which inputs were successfully processed by the primary instance, such that when the primary instance fails and a backup instance is started, the backup store need only be consulted to determine which inputs need to be re-processed by the backup instance before it can assume normal processing.

(32) As described above, when implementing a backup system, the operator needs to consider where to locate, physically and/or logically, that backup system relative to the primary instance and the source of inputs thereto. For example, in the case of physical infrastructure issues or natural disasters which may affect the geographic location of the primary instance, one would not want the backup system to be physically located in the same geographic areas as then it may be vulnerable to the same issues, thereby reducing the effectiveness of the backup system. Accordingly, most operators locate their backup systems in different geographic regions from where the primary

instance is located to minimize the chance that both systems will experience failures due to the same cause. It will be appreciated that many operators deploy multiple backup systems located in disparate geographic regions to further minimize the likelihood of the same event compromising both the primary instance and the availability of at least one backup system.

(33) However, as already stated above, locating the backup system in a geographic location different from where the primary instance is located creates a latency issue with regard to communication of the transactional inputs to both systems to assure that the backup system is synchronized with the primary system should the backup system need to take over, i.e., to minimize transaction reprocessing or loss. One solution may be to receive the transactional inputs from their source at a location that is equidistant, or otherwise subject to substantially equal/symmetric communications latencies, to both the primary instance and the backup system, where the transactions are then relayed to both system from this equidistant location. However, for storage-based backup systems, checkpoints from the primary instance may still need to be periodically communicated between the primary instance and the backup system to minimize the extent to which inputs need to be reprocessed when a failure occurs. Furthermore, in latency sensitive applications, such as financial transaction processing, it may be necessary to locate the primary instance close to the source of transactions sent thereto so as to minimize the operational latency of the application/system. This may necessarily mean the backup system has to be located further from that source of transactions.

(34) As used herein, primary instances of a system and the backup instances/systems, or components thereof, may be geographically/physically and/or logically separated from one another introducing communications latency therebetween, i.e., they may be separated geographically/physically, e.g., located in different physical locations or geographic regions, and/or logically separated, e.g., by one or more interconnecting communications media or other intervening components, such as relays, gateways or switching devices. For example, a communications path of a certain length comprising numerous intervening gateways or switching devices may be characterized by more latency than a longer communications path having fewer such intervening components. More particularly, the distance/length/latency of a given data/communications path interconnecting any two of the described components or other intervening components, whether those components are themselves physically close or not, may introduce latency in the electronic communications therebetween. Further, any asymmetries in the distance/length/latency between the interconnecting data/communications paths, or the number or type of intervening components, whether or not they interconnect the same source and destination end points, may introduce similar asymmetries in the latencies of the electronic communications communicated therethrough.

(35) Further, differences in communications latency of a given communications/network path, or as between two different network paths to a common destination, may be caused by static differences, dynamic differences, or a combination thereof, in the network infrastructure which makes up those network paths, e.g., network switches, wires, wireless connections, etc. Static differences include: media type/characteristics such as cat6 copper cable, fiber optic, satellite, microwave or Wi-Fi; cable length/impedance where a longer and/or higher impedance cable requires a longer time to transit than a shorter and/or lower impedance cable of the same type; number, type and capability of routing and switching devices along the path which impart processing delay to perform their functions; transceivers which transfer/translate messages between different media such as between copper, fiber and wireless media, etc. Generally, static differences are differences which do not change over time, e.g., delays attributable to static characteristics of the network infrastructure. Dynamic differences may include: network load where increased network traffic/congestion may increase latency; interference such as radio frequency interference, sunspots, etc. which may cause errors and retries in the transmission; equipment/media degradation or mechanical issues such as temperature/environmental sensitivity, poor connections or degraded components which may

impart transmission errors or intermittent or varying changes in impedance, capacitive and/or resistive delay, etc. Generally, dynamic latency differences vary over time and may or may not be predictable. Given dynamic latency variations, a network path that has a higher latency as compared to another network path at a particular time may have a lower latency at another time. Dynamic latencies may affect different messages along the same path where, not only may one message transit the network path faster than another message, but one message may overtake another message in transit such as where an earlier forwarded message must be resent by intermediate network component due to an intermittent error and where the earlier message is resent after a later forwarded message is communicated by that intermediate network component. It will be appreciated that static latency differences may be addressed by measuring the latency variances among the different network paths and physically or logically statically compensating for those difference such as by adding an additional length of cable or an electronic fixed delay buffer along a lower latency path to equalize that path to a longer latency path. Alternatively, slower network infrastructure components may be replaced with faster components to decrease latency commensurate with another network path. While some dynamic latency issues may be mitigated using static changes, such as replacing interference-susceptible components with shielded components, implementing proper maintenance and upkeep, etc., it will be appreciated that given the nature of dynamic latencies, such latencies cannot be completely addressed in this manner.

(36) Communications latency differentials/disparities/asymmetries may result in transaction inputs being received at the backup system later than they were received at the primary instance resulting in the backup system operating “behind” the primary instance. That is, at any moment in time where the primary instance is processing a series of transactions, T , $T+1$, $T+2 \dots T+50$, etc., where the primary has processed up to $T+38$, the backup may still be processing transaction $T+2$. Furthermore, at any given moment, transaction inputs, e.g., $T3-37$, may be en route, or otherwise “in flight” or “on the wire” to the backup system and vulnerable to data loss should a failure compromise the mode of communication.

(37) Accordingly, backup systems may be intentionally implemented so as to process transactions “behind” the primary instance, e.g., only processing a given transaction once it is known that the primary instance has already successfully processed that transaction. In a failover situation, the backup instance need only catch up by processing the most recent transaction received and/or attempted by the primary instance and then the backup instance may proceed to take over to replace the failed primary instance to process newly received transactions.

(38) In one implementation, where the components of the electronic trading system are separately deployed and interconnected with a network, a component level backup system may be deployed such as the system using a backup/active copy-cat instance to achieve fault tolerance, see U.S. Pat. No. 7,434,096 ('096), filed on Aug. 11, 2006, entitled “MATCH SERVER FOR A FINANCIAL EXCHANGE HAVING FAULT TOLERANT OPERATION” and U.S. Pat. No. 7,480,827 ('827), filed on Aug. 11, 2006, entitled “FAULT TOLERANCE AND FAILOVER USING ACTIVE COPY-CAT”, assigned to the assignee of the present application, the entirety of all of which are incorporated by reference herein and relied upon.

(39) Both '096 and '827, relate to a fault tolerant failover mechanism allowing the backup instance of a specific component, e.g., the match engine, to take over for the primary instance of that component in a fault situation wherein the primary and backup instances are loosely coupled, i.e., they need not be aware of each other or that they are operating in a fault tolerant environment. As such, the primary instance need not be specifically designed or programmed to interact with the fault tolerant mechanisms. Instead, the primary instance needs only be designed to adhere to specific basic operating guidelines and shut itself down when it cannot do so. By externally controlling the ability of the primary instance to successfully adhere to its operating guidelines, the fault tolerant mechanisms of both '096 and '827 can recognize error conditions and easily failover from the primary instance to the backup instance. In these applications, in contrast with the

disclosed embodiments, a primary instance refers to a single process, thread, or application/component. Therefore, the fault tolerance mechanism only replaces the single application that failed and not the whole system. Further, in these systems, an external mechanism is implemented to force an application into failure, by, for example, preventing a database commit, to deploy the fault tolerant functionality without having to specifically modify the application to incorporate it therein.

(40) As operators seek to enhance the performance of the electronic trading system, one area of focus is on the network connections which are used to interconnect the various components of the system and which may introduce substantial performance degradation, in the form of latency, to the operation of the system. To mitigate such performance issues, the system components may instead be integrated and operated on a single processing platform, i.e., a single- or multi-processor server or tightly coupled/interconnected set of servers, enabling elimination of the higher latency network interconnects and associated latencies introduced thereby in favor of much faster interconnects such as direct processing interconnects, inter-process communication and/or shared memory messaging.

(41) However, the integration of the components into a single processing environment further complicates the provision of fault tolerant operation as the use of component level backup instances would again require, in a failover situation, that the correctly operating components communicate with the backup instance of the failed component over a necessarily higher-latency network/interconnect, necessarily impacting performance. Furthermore, in an integrated implementation any fault which might impact one component is likely to impact all components.

(42) As will be described further below, the disclosed embodiments are implemented in conjunction with an integrated system/architecture that implements the core applications/components/engines of the above described electronic trading system, e.g., the TR, the TP, the RG, and TL, on, for example, a single server computer having at least one processor and memory to allow a technical benefit of using shared memory communication between the core components. The shared memory communication relates to an inter-process communication (IPC) that enables exchanging data between applications/components running at the same time.

(43) Even though the integration of the electronic trading system components enables implementation of all core components on a single server or tightly coupled/interconnected group thereof, this implementation causes unique challenges in how to handle fault tolerance (e.g., starting up and transitioning to a backup) and coordinating state synchronization across core components.

(44) In particular, in one implementation of the disclosed embodiments, in contrast with other systems for implementing fault tolerance, provide fault tolerance to integrated systems which move/integrate core components/applications from separate servers interconnected via a network, each with a separate backup, to a single server/processing environment while making sure all components/applications are backed up and synchronized.

(45) The disclosed embodiments provide a hard backup, i.e., a similar/replicated system implemented on a different server or tightly coupled/integrated set thereof, e.g., located remote from the primary system, to maintain resiliency when the core components/applications are integrated on the same server. Therefore, in contrast with previous fault tolerance systems that replace only failed components, if there are any issues on any of the core components on the primary server, then all of the core components on the primary server fail over to the backup server.

(46) In one implementation, as shown in FIG. 3, the disclosed embodiments enable the provision of at least two or more identical instances of an electronic trading system, e.g., transaction processing servers, where one of those instances serves as the primary instance of the trading system responsible for processing incoming requests for a transaction and the other instances serve as a backup therefore. In particular, each instance is identical to each other but is configured to operate differently depending on its "role", i.e., whether the instance serves as a primary instance or as a

backup. Each instance may include software applications/components/threads executing on one or more processors/processing cores and/or other hardware components, consisting of, or executing on a processing element, server, or the like. Each component is characterized by a current state. Each instance is configured to receive a message via an electronic communications network, process the message to produce a result, and generate a result message including the result which may alter the current state of the instance. While described separately in terms of components and their particular functions/operations, it will be appreciated that the disclosed components of each instance of the disclosed transaction processing server may be implemented as one or more separate applications, computer programs and/or processing threads, which may be independently or dependently executed/implemented, such as substantially in parallel, in an integrated processing environment, such as by one or more physical and/or virtual processors or processing cores of one or more server computers, e.g., tightly coupled/interconnected via low latency interconnects. and all such implementations are contemplated herein. For example, the TR component may be implemented in a networking component of a server or a network switch coupled therewith, i.e., as part of the network interface (the "NIC") to the communications network, so as to receive and process incoming transaction messages from the network as soon as possible. Each component is configured to perform one or more operations including the receipt of the message, processing of the message, and generation of result messages. Each component is further configured to communicate with its counterpart in the other instances via a network/inter-instance network or bus and track the counterpart's current state and what message is being processed on what component via, for example, status messages. To implement the disclosed fault tolerant functionality, each instance is also configured to join a corresponding Fault Tolerance (FT) group. In one embodiment, the Fault Tolerance group may be managed/executed/implemented by a processor, e.g., a fault tolerance processor (FT or FT processor). Instead of creating an instance for each component, each of the components may be configured to initialize and start up in the corresponding single server/processing environment and register/subscribe to a corresponding Fault Tolerance group/FT. (47) In one implementation, one or more of the plurality of instances may be implemented separate from each other in independent environments but in the same physical hardware. In another implementation, one or more of the plurality of instances may be implemented separate from each other and may be configured to be implemented in a plurality of separate physical hardware. Each separate physical hardware may be disposed in different locations, e.g., different positions in the same rack in the same data center, different racks in the same data center, or in separate geographical locations. No matter where they are located, relative to each other, each of the plurality of instances operate in similar manner as described herein.

(48) In one embodiment, each component of a given instance is configured to communicate with another component of that same instance via a shared memory architecture or other inter-process communication or low latency interconnect mechanism.

(49) Each component is configured to start up, initialize communications, and communicate with the corresponding FT via status messages. Each FT is further configured to coordinate and synchronize the state of the components that are running in the same corresponding instance. Once the FT has determined that all corresponding components have started, i.e., once all of the components have joined and registered with a corresponding FT group, i.e., once the FT group has been created, the FT communicates with the other FTs to, for example, vote amongst each other and/or determine or otherwise designate one of the instances to operate in the role of the primary instance and the others to operate in the role of a backup instance. Each FT notifies all components on the corresponding instance whether they are part of a primary or a backup instance. In other words, each component on each corresponding instance knows whether the component belongs to a primary or a backup instance. In one embodiment, each of the FTs determines the primary instance by determining which one of the instances is the first to report that all of the components have been registered with the corresponding FT.

(50) As already indicated, the core components/applications may be grouped in a collection (or Fault Tolerance Group). In one embodiment, each collection is implemented in the same server/instance. The whole group of core components is replicated in the backup groups/instance(s)/server(s). Each Fault Tolerance group may have N number of applications, for example FIG. 3 shows 3, but it could be 4 (as shown in FIG. 5) or any other amount. When the core components/applications are started, all of the components within the same group or instance, need to have the same role, either a primary role or a backup role. As mentioned above, the FT is configured to get the role assigned to each corresponding group/instance/server.

(51) In implementations where the functions of one or more of the disclosed components are implemented as an integrated or otherwise common component/thread/application, those integrated components may collectively register with the FT as described.

(52) In one implementation, each collection of core components are separate from each other and are configured to run on separate physical hardware.

(53) In one implementation, one or more of the core components may be implemented by one or more processors. In another implementation, one or more of the plurality of components may be implemented separate from each other in independent environments but in the same physical hardware. In another implementation, one or more of the plurality of components may be implemented separate from each other in different physical hardware. In one implementation, an FPGA based implementation may be used that permits the components of each instance to be collocated, such as via a custom interface, or otherwise closely interconnected with networking equipment, such as a router or switch, e.g. via a backplane thereof. This would allow each component to communicate as quickly as possible and avoid the latency introduced, not only by having to route the messages over conventional networking media, but also by the communications protocols, e.g. Transport Control Protocol ("TCP"), used to perform that routing.

(54) While the disclosed embodiments will be discussed with respect to a single backup instance of the electronic trading system, it will be appreciated that more than one backup instance may be implemented to provide further redundancy. In one implementation, the disclosed embodiments implement multiple backup instances, which may be geographically dispersed, wherein the primary instance interacts with each of the backup instances in an identical manner as will be described.

(55) The disclosed embodiments may provide order message information to external recipients via two separate modes of network communications: a direct/private path used by a customer entering an order and receiving a confirmation of receipt and/or fill thereof, e.g., a transaction message (TM) published to the customer, such as via iLink, an order routing and communication protocol provided by the Chicago Mercantile Exchange Inc., and a public path, e.g., the trade summary message published to all market participants, e.g., via multicast. The TM interface facilitates the entry, modification, and cancellation of orders, as well as receipt of order confirmation and fill information.

(56) In one implementation, one of the differences between the operation of an instance as a primary or as a backup is communication with external recipients, e.g., the customer and all market participants. In particular, only the primary instance transmits the results to external recipients via the private and public path. Results produced by the backup instance are not transmitted to external recipients, e.g., the results are withheld by the backup instance or otherwise prevented from being transmitted to or otherwise reaching the external recipients.

(57) In one implementation, as the primary instance processes incoming transactions, the primary instance transmits the current state of each of its components to the backup instance, i.e., to the corresponding components thereof. As will be described, because the primary instance first receives and initiates processing of incoming transactions, the state of each of the components of the backup instance will be "behind" the corresponding primary instance component, i.e., the state of each of the components of the backup instance will necessarily correspond to the state of each of the components of the primary instance subsequent to the processing of a message preceding the

message currently being processed by the primary instance. Specifically, as discussed further below, the backup instance, as a byproduct of its dependency on the primary instance, will necessarily process incoming order requests/messages “behind” the primary instance, e.g., each component thereof only performing its processing task of a given incoming order message once it is known that the corresponding component of the primary instance has already successfully processed that message. In a failover situation, as will be described, the backup instance, i.e., each component thereof, needs only catch up by processing the most recent transaction attempted by the primary instance and then the backup instance may proceed to take over to replace the failed primary instance.

(58) In particular, each of the components of the primary instance is configured to execute an operation in the furtherance of processing an incoming order message and to communicate a result message, e.g., an electronic result message, that includes the result of the operation to the corresponding component of each backup instance. Subsequent to the receipt of the result message, each corresponding component of the backup instance processes the given message and compares a result of the operation thereof with the result indicated by the received result message to determine whether one of the components of the primary instance or the corresponding component of the backup instance is operating correctly. In other words, each component of the backup instance does not start processing a given message until it has received the corresponding result message from the primary instance.

(59) The disclosed embodiments further relate to a fault tolerance system/process which provides redundancy and error/failure recovery to an integrated electronic trading system as described. Specifically, when the disclosed fault tolerance system determines that a failure has occurred on the primary instance, the fault tolerance system designates one of the backup instances as a new primary instance. The backup instance designated to become the new primary instance is notified that it needs to become the primary instance and, depending on the type of failure, the primary instance may continue to queue messages until all components in the designated backup instance have caught up processing queued messages. Therefore, all messages are processed in the correct sequence and all messages are processed before failing over to the designated backup instance. Depending on the type of failure, when the FT coupled with the primary instance determines that the failure has occurred in the primary instance, each of the FTs is configured to cause the primary instance to at least de-couple from the instances and/or from the electronic communications network.

(60) In particular, each of the FTs is configured to determine whether a fault occurred in the primary instance. When a FT determines that a failure occurred in the primary instance, the FT is configured to notify the other FTs that a failure of the primary instance has occurred. The FTs then vote to designate one of the backup instances as a new primary instance. In particular, the FTs cause one of the backup instances to be designated as a new primary instance. The backup instance designated as the new primary instance is configured to synchronize the current state of each of the plurality of components thereof based on received electronic messages indicative of a result of the operation of the corresponding component of the primary instance, and process any incoming requests not processed by the primary instance before commencing processing of new incoming requests.

(61) In one implementation, each of components of each instance is coupled with the corresponding FT such that the corresponding FT is able to detect a failure of any of the components of the instance. A failure on any of the components of the primary instance is deemed to indicate a failure of the primary instance.

(62) In one implementation, each FT periodically receives a communication, e.g., a heartbeat message, from the others FTs. A failure to receive the communication from one of the FTs is deemed to indicate a failure of the instance coupled therewith. The failure to receive the communication is indicative that there may be a network failure.

(63) As already noted, the disclosed embodiments provide an improved transaction processing system that enables fault tolerant operation of integrated electronic trading systems which eliminate the network interconnects and associated latencies between components introduced thereby in favor of much faster interconnects such as inter-process communication and shared memory messaging.

(64) The disclosed embodiments further provide an improved system resiliency and disaster recovery mechanism for a deterministic transaction processing system which can adaptively accommodate processing and communications latencies, for example, in a cloud/shared resource deployment, to mitigate data loss and performance impact while minimizing failure recovery costs.

(65) The provision of a fault tolerance system for a transaction processing system which minimize, if not eliminate, both the amount of data, i.e., the number of transactions, which can be lost during a failure and the performance impact on the system is a technical implementation and problems therewith, such as data loss, processing and communications latency variation, processing performance degradation, are technical problems which can affect computer processes and systems which rely on electronic message based communication for operation. As such, the disclosed embodiments provide technical solutions to these technical problems.

(66) The disclosed embodiments solve problems which uniquely arise in the fields of computer technology and electronic communication. The disclosed embodiments are rooted in computer technology in order to overcome problems specifically arising in computer systems and inter-process communications. Indeed, the subject technology improves the functioning of the computer by, for example, minimizing, if not eliminate, both the amount of data, i.e., the number of transactions, which can be lost during a failure, the performance impact on the system and the costs to recover in the event of failure.

(67) Therefore, the disclosed embodiments provide a specific and practical application which improves upon prior messaging and electronic trading platforms and provides additional functionality not previously provided.

(68) When a component, device, element, or the like of the present disclosure is described as having a purpose or performing an operation, function, or the like, the component, device, or element should be considered herein as being “configured to” or “operative to” meet that purpose or to perform that operation or function.

(69) Computing Environment

(70) The embodiments may be described in terms of a computing system. The particular examples identify a specific set of components useful in a futures and options exchange. However, many of the components and inventive features are readily adapted to other electronic trading environments. The specific examples described herein may teach specific protocols and/or interfaces, although it should be understood that the principles involved may be extended to, or applied in, other protocols and interfaces.

(71) It should be appreciated that the plurality of entities utilizing or involved with the disclosed embodiments, e.g., the market participants, may be referred to by other nomenclature reflecting the role that the particular entity is performing with respect to the disclosed embodiments and that a given entity may perform more than one role depending upon the implementation and the nature of the particular transaction being undertaken, as well as the entity's contractual and/or legal relationship with another market participant and/or the exchange.

(72) A trading network environment for implementing trading systems and methods is shown in FIG. 1. FIG. 1 shows a block diagram of an electronic trading system/architecture **100** which may be implemented as an instance of the integrated electronic trading system operating on a server as described in more detail herein. The system **100** receives messages that include orders and transmits market data related to orders and trades to users, such as via a wide area network **162** and/or a local area network **160** and computer devices **150**, **152**, **154**, **156** and **158**, as described herein, coupled with the exchange computer system **100**.

(73) Herein, the phrase “coupled with” is defined to mean directly connected to or indirectly connected through one or more intermediate components. Such intermediate components may include both hardware and software based components. Further, to clarify the use in the pending claims and to hereby provide notice to the public, the phrases “at least one of <A>, , . . . and <N>” or “at least one of <A>, , . . . <N>, or combinations thereof” are defined by the Applicant in the broadest sense, superseding any other implied definitions herebefore or hereinafter unless expressly asserted by the Applicant to the contrary, to mean one or more elements selected from the group comprising A, B, . . . and N, that is to say, any combination of one or more of the elements A, B, . . . or N including any one element alone or in combination with one or more of the other elements which may also include, in combination, additional elements not listed.

(74) The system **100** may be implemented with one or more mainframe, desktop, or other computers, such as the example computer **200** described herein with respect to FIG. 2. A user database **102** may be provided which includes information identifying traders and other users of the exchange computer system **100**, such as account numbers or identifiers, usernames, and passwords. An account data module **104** may be provided which may process account information that may be used during trades.

(75) A match engine module **106** may be included to match bid and offer prices and may be implemented with software that executes one or more algorithms for matching bids and offers. A trade database **108** may be included to store information identifying trades and descriptions of trades. In particular, trade database **108** may store information identifying the time that a trade took place and the contract price.

(76) An order book module **110** may be included to compute or otherwise determine current bid and offer prices, e.g., in a continuous auction market, or also operate as an order accumulation buffer for a batch auction market.

(77) A market data module **112** may be included to collect market data and prepare the data for transmission to users. For example, the market data module **112** may prepare the market data feeds described herein.

(78) A risk management module **114** may be included to compute and determine a user's risk utilization in relation to the user's defined risk thresholds. The risk management module **114** may also be configured to determine risk assessments or exposure levels in connection with positions held by a market participant. The risk management module **114** may be configured to administer, manage, or maintain one or more margining mechanisms implemented by the exchange computer system **100**. Such administration, management or maintenance may include managing a number of database records reflective of margin accounts of the market participants.

(79) A message management module **116** may be included to, among other things, receive, and extract orders from, electronic data transaction request messages. The message management module **116** may define a point of ingress into the system **100** where messages are ordered and considered to be received by the system **100**. This may be considered a point of determinism in the system **100** that defines the earliest point where the system **100** can ascribe an order of receipt to arriving messages. The point of determinism may or may not be at or near the demarcation point between the system **100** and a public/internet network infrastructure. The message management module **116** processes messages by interpreting the contents of a message based on the message transmit protocol, such as the transmission control protocol (“TCP”), to provide the content of the message for further processing by the exchange computer system.

(80) The message management module **116** may also be configured to detect characteristics of an order for a transaction to be undertaken in an electronic marketplace. For example, the message management module **116** may identify and extract order content such as a price, product, volume, and associated market participant for an order. The message management module **116** may also identify and extract data indicating an action to be executed by the system **100** with respect to the extracted order. For example, the message management module **116** may determine the transaction

type of the transaction requested in a given message. A message may include an instruction to perform a type of transaction. The transaction type may be, in one embodiment, a request/offer/order to either buy or sell a specified quantity or units of a financial instrument at a specified price or value. The message management module **116** may also identify and extract other order information and other actions associated with the extracted order. All extracted order characteristics, other information, and associated actions extracted from a message for an order may be collectively considered an order as described and referenced herein.

(81) Order or message characteristics may include, for example, the state of the system after a message is received, arrival time (e.g., the time a message arrives at the TR that is the point of ingress/entry and/or egress/departure for all transactions, i.e., the network traffic/packets containing the data therefore), message type (e.g., new, modify, cancel), and the number of matches generated by a message. Order or message characteristics may also include market participant side (e.g., buyer or seller) or time in force (e.g., a good until end of day order that is good for the full trading day, a good until canceled order that rests on the order book until matched, or a fill or kill order that is canceled if not filled immediately, or a fill and kill order (FOK) that is filled to the maximum amount possible based on the state of the order book at the time the FOK order is processed, and any remaining or unfilled/unsatisfied quantity is not stored on the books or allowed to rest).

(82) An order processing module **118** may be included to decompose delta-based, spread instrument, bulk, and other types of composite orders for processing by the order book module **110** and/or the match engine module **106**. The order processing module **118** may also be used to implement one or more procedures related to clearing an order. The order may be communicated from the message management module **116** to the order processing module **118**. The order processing module **118** may be configured to interpret the communicated order, and manage the order characteristics, other information, and associated actions as they are processed through an order book module **110** and eventually transacted on an electronic market. For example, the order processing module **118** may store the order characteristics and other content and execute the associated actions. In an embodiment, the order processing module **118** may execute an associated action of placing the order into an order book for an electronic trading system managed by the order book module **110**. In an embodiment, placing an order into an order book and/or into an electronic trading system may be considered a primary action for an order. The order processing module **118** may be configured in various arrangements and may be configured as part of the order book module **110**, part of the message management module **116**, or as an independent functioning module.

(83) As an intermediary to electronic trading transactions, the exchange bears a certain amount of risk in each transaction that takes place. To that end, the clearing house implements risk management mechanisms to protect the exchange. One or more of the modules of the system **100** may be configured to determine settlement prices for constituent contracts, such as deferred month contracts, of spread instruments, such as for example, settlement module **120**. A settlement module **120** (or settlement processor or other payment processor) may be included to provide one or more functions related to settling or otherwise administering transactions cleared by the exchange. Settlement module **120** of the system **100** may implement one or more settlement price determination techniques. Settlement-related functions need not be limited to actions or events occurring at the end of a contract term. For instance, in some embodiments, settlement-related functions may include or involve daily or other mark to market settlements for margining purposes. In some cases, the settlement module **120** may be configured to communicate with the trade database **108** (or the memory(ies) on which the trade database **108** is stored) and/or to determine a payment amount based on a spot price, the price of the futures contract or other financial instrument, or other price data, at various times. The determination may be made at one or more points in time during the term of the financial instrument in connection with a margining mechanism. For example, the settlement module **120** may be used to determine a mark to market

amount on a daily basis during the term of the financial instrument. Such determinations may also be made on a settlement date for the financial instrument for the purposes of final settlement.

(84) In some embodiments, the settlement module **120** may be integrated to any desired extent with one or more of the other modules or processors of the system **100**. For example, the settlement module **120** and the risk management module **114** may be integrated to any desired extent. In some cases, one or more margining procedures or other aspects of the margining mechanism(s) may be implemented by the settlement module **120**.

(85) A fault tolerance module **124**, shown in more detail in FIGS. **12A(1)**, **12A(2)**, **12A(3)**, and **12B** described in more detail below, may be included to, using the disclosed embodiments, implement, facilitate, and manage both the reliable transmission of copies of messages/state of the components of the system **100** and, in the event of a failure, the recovery thereof, e.g., among other modules or components of the system **100**, as described herein. This fault tolerance module **124** may be part of, or include, a messaging module, not shown, which implements inter-component message communication.

(86) One or more of the above-described modules of the system **100** may be used to gather or obtain data to support the settlement price determination, as well as a subsequent margin requirement determination. For example, the order book module **110** and/or the market data module **112** may be used to receive, access, or otherwise obtain market data, such as bid-offer values of orders currently on the order books. The trade database **108** may be used to receive, access, or otherwise obtain trade data indicative of the prices and volumes of trades that were recently executed in a number of markets. In some cases, transaction data (and/or bid/ask data) may be gathered or obtained from open outcry pits and/or other sources and incorporated into the trade and market data from the electronic trading system(s). It should be appreciated that concurrent processing limits may be defined by or imposed separately or in combination on one or more of the trading system components.

(87) The disclosed mechanisms may be implemented at any logical and/or physical point(s), or combinations thereof, at which the relevant information/data (e.g., message traffic and responses thereto) may be monitored or flows or is otherwise accessible or measurable, including one or more gateway devices, modems, the computers, or terminals of one or more market participants, e.g., client computers, etc.

(88) One skilled in the art will appreciate that one or more modules described herein may be implemented using, among other things, a tangible computer-readable medium comprising computer-executable instructions (e.g., executable software code). Alternatively, modules may be implemented as software code, firmware code, specifically configured hardware or processors, and/or a combination of the aforementioned. For example, the modules may be embodied as part of an exchange **100** for financial instruments. It should be appreciated the disclosed embodiments may be implemented as a different or separate module of the exchange computer system **100**, or a separate computer system coupled with the exchange computer system **100** so as to have access to margin account record, pricing, and/or other data. As described herein, the disclosed embodiments may be implemented as a centrally accessible system or as a distributed system, e.g., where some of the disclosed functions are performed by the computer systems of the market participants.

(89) The trading network environment shown in FIG. **1** includes exemplary computer devices **150**, **152**, **154**, **156** and **158** which depict different exemplary methods or media by which a computer device may be coupled with the system **100** or by which a user may communicate, e.g., send and receive, trade, clearing or other information therewith, such as by using messaging via the messaging module **122**. It should be appreciated that the types of computer devices deployed by traders and the methods and media by which they communicate with the system **100** is implementation dependent and may vary and that not all of the depicted computer devices and/or means/media of communication may be used and that other computer devices and/or means/media of communications, now available or later developed may be used. Each computer device, which

may comprise a computer **200** described in more detail with respect to FIG. 2, may include a central processor, specifically configured or otherwise, that controls the overall operation of the computer and a system bus that connects the central processor to one or more conventional components, such as a network card or modem. Each computer device may also include a variety of interface units and drives for reading and writing data or files and communicating with other computer devices and with the system **100**. Depending on the type of computer device, a user can interact with the computer with a keyboard, pointing device, microphone, pen device or other input device now available or later developed.

(90) An exemplary computer device **150** is shown directly connected to system **100**, such as via a TI line, a common local area network (LAN) or other wired and/or wireless medium for connecting computer devices, such as the network **220** shown in FIG. 2 and described with respect thereto. The exemplary computer device **150** is further shown connected to a radio **168**. The user of radio **168**, which may include a cellular telephone, smart phone, or other wireless proprietary and/or non-proprietary device, may be a trader or exchange employee. The radio user may transmit orders or other information to the exemplary computer device **150** or a user thereof. The user of the exemplary computer device **150**, or the exemplary computer device **150** alone and/or autonomously, may then transmit the trade or other information to the exchange computer system **100**.

(91) Exemplary computer devices **152** and **154** are coupled with a local area network (“LAN”) **160** which may be configured in one or more of the well-known LAN topologies, e.g., star, daisy chain, etc., and may use a variety of different protocols, such as Ethernet, TCP/IP, etc. The exemplary computer devices **152** and **154** may communicate with each other and with other computers and other devices which are coupled with the LAN **160**. Computers and other devices may be coupled with the LAN **160** via twisted pair wires, coaxial cable, fiber optics or other wired or wireless media. As shown in FIG. 1, an exemplary wireless personal digital assistant device (“PDA”) **158**, such as a mobile telephone, tablet based compute device, or other wireless device, may communicate with the LAN **160** and/or the Internet **162** via radio waves, such as via WiFi, Bluetooth and/or a cellular telephone based data communications protocol. PDA **158** may also communicate with exchange computer system **100** via a conventional wireless hub **164**.

(92) FIG. 1 also shows the LAN **160** coupled with a wide area network (“WAN”) **162** which may be comprised of one or more public or private wired or wireless networks. In one embodiment, the WAN **162** includes the Internet **162**. The LAN **160** may include a router to connect LAN **160** to the Internet **162**. Exemplary computer device **156** is shown coupled directly to the Internet **162**, such as via a modem, DSL line, satellite dish or any other device for connecting a computer device to the Internet **162** via a service provider therefore as is known. LAN **160** and/or WAN **162** may be the same as the network **220** shown in FIG. 2 and described with respect thereto.

(93) Users of the system **100** may include one or more market makers **166** which may maintain a market by providing constant bid and offer prices for a derivative or security to the system **100**, such as via one of the exemplary computer devices depicted. The system **100** may also exchange information with other match or trade engines, such as trade engine **170**. One skilled in the art will appreciate that numerous additional computers and systems may be coupled to the system **100**. Such computers and systems may include clearing, regulatory and fee systems.

(94) The operations of computer devices and systems shown in FIG. 1 may be controlled by computer-executable instructions stored on a non-transitory computer-readable medium. For example, the exemplary computer device **152** may store computer-executable instructions for receiving order information from a user, transmitting that order information to system **100** in electronic messages, extracting the order information from the electronic messages, executing actions relating to the messages, and/or calculating values from characteristics of the extracted order to facilitate matching orders and executing trades. In another example, the exemplary computer device **154** may include computer-executable instructions for receiving market data from

exchange computer system **100** and displaying that information to a user.

(95) Numerous additional servers, computers, handheld devices, personal digital assistants, telephones, and other devices may also be connected to the system **100**. Moreover, one skilled in the art will appreciate that the topology shown in FIG. **1** is merely an example and that the components shown in FIG. **1** may include other components not shown and be connected by numerous alternative topologies.

(96) Referring now to FIG. **2**, an illustrative embodiment of a general computer system **200** is shown. The computer system **200** can include a set of instructions that can be executed to cause the computer system **200** to perform any one or more of the methods or computer based functions disclosed herein. The computer system **200** may operate as a standalone device or may be connected, e.g., using a network, to other computer systems or peripheral devices. Any of the components discussed herein, such as processor **202**, may be a computer system **200** or a component in the computer system **200**. The computer system **200** may be specifically configured to implement a match engine, margin processing, payment or clearing function on behalf of an exchange, such as the Chicago Mercantile Exchange, of which the disclosed embodiments are a component thereof.

(97) In a networked deployment, the computer system **200** may operate in the capacity of a server or as a client user computer in a client-server user network environment, or as a peer computer system in a peer-to-peer (or distributed) network environment. The computer system **200** can also be implemented as or incorporated into various devices, such as a personal computer (PC), a tablet PC, a set-top box (STB), a personal digital assistant (PDA), a mobile device, a palmtop computer, a laptop computer, a desktop computer, a communications device, a wireless telephone, a land-line telephone, a control system, a camera, a scanner, a facsimile machine, a printer, a pager, a personal trusted device, a web appliance, a network router, switch or bridge, or any other machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. In a particular embodiment, the computer system **200** can be implemented using electronic devices that provide voice, video, or data communication. Further, while a single computer system **200** is illustrated, the term “system” shall also be taken to include any collection of systems or sub-systems that individually or jointly execute a set, or multiple sets, of instructions to perform one or more computer functions.

(98) As illustrated in FIG. **2**, the computer system **200** may include a processor **202**, e.g., a central processing unit (CPU), a graphics processing unit (GPU), or both. The processor **202** may be a component in a variety of systems. For example, the processor **202** may be part of a standard personal computer or a workstation. The processor **202** may be one or more general processors, digital signal processors, specifically configured processors, application specific integrated circuits, field programmable gate arrays, servers, networks, digital circuits, analog circuits, combinations thereof, or other now known or later developed devices for analyzing and processing data. The processor **202** may implement a software program, such as code generated manually (i.e., programmed).

(99) The computer system **200** may include a memory **204** that can communicate via a bus **208**. The memory **204** may be a main memory, a static memory, or a dynamic memory. The memory **204** may include, but is not limited to, computer readable storage media such as various types of volatile and non-volatile storage media, including but not limited to random access memory, read-only memory, programmable read-only memory, electrically programmable read-only memory, electrically erasable read-only memory, flash memory, magnetic tape or disk, optical media and the like. In one embodiment, the memory **204** includes a cache or random access memory for the processor **202**. In alternative embodiments, the memory **204** is separate from the processor **202**, such as a cache memory of a processor, the system memory, or other memory. The memory **204** may be an external storage device or database for storing data. Examples include a hard drive, compact disc (“CD”), digital video disc (“DVD”), memory card, memory stick, floppy disc,

universal serial bus (“USB”) memory device, or any other device operative to store data. The memory **204** is operable to store instructions executable by the processor **202**. The functions, acts or tasks illustrated in the figures or described herein may be performed by the programmed processor **202** executing the instructions **212** stored in the memory **204**. The functions, acts or tasks are independent of the particular type of instructions set, storage media, processor or processing strategy and may be performed by software, hardware, integrated circuits, firmware, micro-code, and the like, operating alone or in combination. Likewise, processing strategies may include multiprocessing, multitasking, parallel processing, and the like.

(100) As shown, the computer system **200** may further include a display unit **214**, such as a liquid crystal display (LCD), an organic light emitting diode (OLED), a flat panel display, a solid state display, a cathode ray tube (CRT), a projector, a printer or other now known or later developed display device for outputting determined information. The display **214** may act as an interface for the user to see the functioning of the processor **202**, or specifically as an interface with the software stored in the memory **204** or in the drive unit **206**.

(101) Additionally, the computer system **200** may include an input device **216** configured to allow a user to interact with any of the components of system **200**. The input device **216** may be a number pad, a keyboard, or a cursor control device, such as a mouse, or a joystick, touch screen display, remote control, or any other device operative to interact with the system **200**.

(102) In a particular embodiment, as depicted in FIG. **2**, the computer system **200** may also include a disk or optical drive unit **206**. The disk drive unit **206** may include a computer-readable medium **210** in which one or more sets of instructions **212**, e.g., software, can be embedded. Further, the instructions **212** may embody one or more of the methods or logic as described herein. In a particular embodiment, the instructions **212** may reside completely, or at least partially, within the memory **204** and/or within the processor **202** during execution by the computer system **200**. The memory **204** and the processor **202** also may include computer-readable media as discussed herein.

(103) The present disclosure contemplates a computer-readable medium that includes instructions **212** or receives and executes instructions **212** responsive to a propagated signal, so that a device connected to a network **220** can communicate voice, video, audio, images, or any other data over the network **220**. Further, the instructions **212** may be transmitted or received over the network **220** via a communication interface **218**. The communication interface **218** may be a part of the processor **202** or may be a separate component. The communication interface **218** may be created in software or may be a physical connection in hardware. The communication interface **218** is configured to connect with a network **220**, external media, the display **214**, or any other components in system **200**, or combinations thereof. The connection with the network **220** may be a physical connection, such as a wired Ethernet connection or may be established wirelessly. Likewise, the additional connections with other components of the system **200** may be physical connections or may be established wirelessly.

(104) The network **220** may include wired networks, wireless networks, or combinations thereof. The wireless network may be a cellular telephone network, an 802.11, 802.16, 802.20, or WiMax network. Further, the network **220** may be a public network, such as the Internet, a private network, such as an intranet, or combinations thereof, and may utilize a variety of networking protocols now available or later developed including, but not limited to, TCP/IP based networking protocols.

(105) Embodiments of the subject matter and the functional operations described in this specification can be implemented in digital electronic circuitry, or in computer software, firmware, or hardware, including the structures disclosed in this specification and their structural equivalents, or in combinations of one or more of them. Embodiments of the subject matter described in this specification can be implemented as one or more computer program products, i.e., one or more modules of computer program instructions encoded on a computer readable medium for execution by, or to control the operation of, data processing apparatus. While the computer-readable medium is shown to be a single medium, the term “computer-readable medium” includes a single medium

or multiple media, such as a centralized or distributed database, and/or associated caches and servers that store one or more sets of instructions. The term “computer-readable medium” shall also include any medium that is capable of storing, encoding, or carrying a set of instructions for execution by a processor or that cause a computer system to perform any one or more of the methods or operations disclosed herein. The computer readable medium can be a machine-readable storage device, a machine-readable storage substrate, a memory device, or a combination of one or more of them. The term “data processing apparatus” encompasses all apparatus, devices, and machines for processing data, including by way of example a programmable processor, a computer, or multiple processors or computers. The apparatus can include, in addition to hardware, code that creates an execution environment for the computer program in question, e.g., code that constitutes processor firmware, a protocol stack, a database management system, an operating system, or a combination of one or more of them.

(106) In a particular non-limiting, exemplary embodiment, the computer-readable medium can include a solid-state memory such as a memory card or other package that houses one or more non-volatile read-only memories. Further, the computer-readable medium can be a random-access memory or other volatile re-writable memory. Additionally, the computer-readable medium can include a magneto-optical or optical medium, such as a disk or tapes or other storage device to capture carrier wave signals such as a signal communicated over a transmission medium. A digital file attachment to an e-mail or other self-contained information archive or set of archives may be considered a distribution medium that is a tangible storage medium. Accordingly, the disclosure is considered to include any one or more of a computer-readable medium or a distribution medium and other equivalents and successor media, in which data or instructions may be stored.

(107) In an alternative embodiment, dedicated or otherwise specifically configured hardware implementations, such as application specific integrated circuits, programmable logic arrays and other hardware devices, can be constructed to implement one or more of the methods described herein. Applications that may include the apparatus and systems of various embodiments can broadly include a variety of electronic and computer systems. One or more embodiments described herein may implement functions using two or more specific interconnected hardware modules or devices with related control and data signals that can be communicated between and through the modules, or as portions of an application-specific integrated circuit. Accordingly, the present system encompasses software, firmware, and hardware implementations.

(108) In accordance with various embodiments of the present disclosure, the methods described herein may be implemented by software programs executable by a computer system. Further, in an exemplary, non-limited embodiment, implementations can include shared or proprietary distributed processing, e.g., a cloud/shared host-based implementation, component/object distributed processing, and parallel processing. Alternatively, virtual computer system processing, e.g., using proprietary and/or cloud/shared host-based implementations, can be constructed to implement one or more of the methods or functionality as described herein.

(109) Although the present specification describes components and functions that may be implemented in particular embodiments with reference to particular standards and protocols, the invention is not limited to such standards and protocols. For example, standards for Internet and other packet switched network transmission (e.g., TCP/IP, UDP/IP, HTML, HTTP, HTTPS) represent examples of the state of the art. Such standards are periodically superseded by faster or more efficient equivalents having essentially the same functions. Accordingly, replacement standards and protocols having the same or similar functions as those disclosed herein are considered equivalents thereof.

(110) A computer program (also known as a program, software, software application, script, or code) can be written in any form of programming language, including compiled or interpreted languages, and it can be deployed in any form, including as a standalone program or as a module, component, subroutine, or other unit suitable for use in a computing environment. A computer

program does not necessarily correspond to a file in a file system. A program can be stored in a portion of a file that holds other programs or data (e.g., one or more scripts stored in a markup language document), in a single file dedicated to the program in question, or in multiple coordinated files (e.g., files that store one or more modules, sub programs, or portions of code). A computer program can be deployed to be executed on one computer or on multiple computers that are located at one site or distributed across multiple sites and interconnected by a communication network.

(111) The processes and logic flows described in this specification can be performed by one or more programmable processors executing one or more computer programs to perform functions by operating on input data and generating output. The processes and logic flows can also be performed by, and apparatus can also be implemented as, special purpose logic circuitry, e.g., an FPGA (field programmable gate array) or an ASIC (application specific integrated circuit).

(112) Processors suitable for the execution of a computer program include, by way of example, special purpose microprocessors, and anyone or more processors of any kind of digital computer. Generally, a processor will receive instructions and data from a read only memory or a random-access memory or both. The essential elements of a computer are a processor for performing instructions and one or more memory devices for storing instructions and data. Generally, a computer will also include, or be operatively coupled to receive data from or transfer data to, or both, one or more mass storage devices for storing data, e.g., magnetic, magneto optical disks, or optical disks. However, a computer need not have such devices. Moreover, a computer can be embedded in another device, e.g., a mobile telephone, a personal digital assistant (PDA), a mobile audio player, a Global Positioning System (GPS) receiver, to name just a few. Computer readable media suitable for storing computer program instructions and data include all forms of non-volatile memory, media, and memory devices, including by way of example semiconductor memory devices, e.g., EPROM, EEPROM, and flash memory devices; magnetic disks, e.g., internal hard disks or removable disks; magneto optical disks; and CD ROM and DVD-ROM disks. The processor and the memory can be supplemented by, or incorporated in, special purpose logic circuitry.

(113) As used herein, the terms “microprocessor” or “general-purpose processor” (“GPP”) may refer to a hardware device that fetches instructions and data from a memory or storage device and executes those instructions (for example, an Intel Xeon processor or an AMD Opteron processor) to then, for example, process the data in accordance therewith. The term “reconfigurable logic” may refer to any logic technology whose form and function can be significantly altered (i.e., reconfigured) in the field post-manufacture as opposed to a microprocessor, whose function can change post-manufacture, e.g., via computer executable software code, but whose form, e.g., the arrangement/layout and interconnection of logical structures, is fixed at manufacture. The term “software” may refer to data processing functionality that is deployed on a GPP. The term “firmware” may refer to data processing functionality that is deployed on reconfigurable logic. One example of a reconfigurable logic is a field programmable gate array (“FPGA”) which is a reconfigurable integrated circuit. An FPGA may contain programmable logic components called “logic blocks,” and a hierarchy of reconfigurable interconnects that allow the blocks to be “wired together,” somewhat like many (changeable) logic gates that can be inter-wired in (many) different configurations. Logic blocks may be configured to perform complex combinatorial functions, or merely simple logic gates like AND, OR, NOT and XOR. An FPGA may further include memory elements, which may be simple flip-flops or more complete blocks of memory.

(114) To provide for interaction with a user, embodiments of the subject matter described in this specification can be implemented on a device having a display, e.g., a CRT (cathode ray tube) or LCD (liquid crystal display) monitor, for displaying information to the user and a keyboard and a pointing device, e.g., a mouse or a trackball, by which the user can provide input to the computer. Other kinds of devices can be used to provide for interaction with a user as well. Feedback

provided to the user can be any form of sensory feedback, e.g., visual feedback, auditory feedback, or tactile feedback. Input from the user can be received in any form, including acoustic, speech, or tactile input.

(115) Embodiments of the subject matter described in this specification can be implemented in a computing system that includes a back end component, e.g., a data server, or that includes a middleware component, e.g., an application server, or that includes a front end component, e.g., a client computer having a graphical user interface or a Web browser through which a user can interact with an implementation of the subject matter described in this specification, or any combination of one or more such back end, middleware, or front end components. The components of the system can be interconnected by any form or medium of digital data communication, e.g., a communication network. Examples of communication networks include a local area network (“LAN”) and a wide area network (“WAN”), e.g., the Internet.

(116) The computing system can include clients and servers. A client and server are generally remote from each other and typically interact through a communication network. The relationship of client and server arises by virtue of computer programs running on the respective computers and having a client-server relationship with each other.

(117) Component Integrated System

(118) FIG. 3 depicts a block diagram of a system **300** for implementing an integrated component and fault tolerant system according to one embodiment.

(119) The system **300** includes a plurality of transaction processing servers **306** coupled with each other via an electronic communications network such as the network **220** shown in FIG. 2. Each of the plurality of transaction processing servers **306** is operative to receive an incoming request for a transaction, process the received request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and previously received request and generate a result indicative thereof. One of the plurality of transaction processing servers **306** is designated as a primary instance **306a** and each of the other transaction processing servers **306b** and **306c** is designated as a backup instance. Each of other transaction servers **306** designated as the backup instance **306b** is configured to receive a copy of the incoming request from the transaction server designated **306a** as the primary instance. The transaction processing server **306a** designated as the primary instance at least commences processing of the incoming request prior to the transaction processing servers **306b** and **306c** designated as backup instances.

(120) In one embodiment, only the transaction processing server **306a** designated as the primary instance is configured to transmit a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and previously received request to a recipient external to the system.

(121) In one embodiment, one or more of the plurality of transaction processing servers **306** is located in a geographic region different from a location of one or more others of the plurality of transaction processing servers **306**.

(122) FIG. 4 depicts a detailed block diagram of a transaction processing server **306**. Referring to FIGS. 3 and 4, each of the plurality of transaction processing servers **306** may include a plurality of components **314**. Each component **314** may be implemented as a separate component or as one or more logic components, such as part of one more of the modules of the exchange computing system **100** described above, such as on an FPGA that may include a memory or reconfigurable component to store logic and a processing component to execute the stored logic, or as computer program logic, stored in a memory **304**, or other non-transitory computer readable medium, and executable by a processor **302**, such as the processor **202** and the memory **204** described above with respect to FIG. 2.

(123) Each of the plurality of components **314** in each of the corresponding transaction processing server **306** is coupled with each other and is characterized by a current state. Each of the plurality of components **314** is configured to perform one or more of receiving an electronic message,

processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof. In one embodiment, each of the plurality of components **314** are coupled with each other via a shared memory architecture, e.g., where messages may be exchanged via one component writing data indicative of the message to one or more common memory locations and the other component reading therefrom.

(124) The system **300** further includes a plurality of fault tolerance processors **316** such as the processor **202** described above with respect to FIG. 2. As detailed above, the plurality of fault tolerance processors/processing components **316** are configured to execute/implement the FTs mentioned above. The system **300** further includes a fault tolerance memory (not shown) such as the memory **204** described above with respect to FIG. 2. Each of the plurality of fault tolerance processors **316** may be coupled with a fault tolerant memory (not shown) storing computer executable instructions executable by the fault tolerance processor **316** to perform the operations described below.

(125) Each of the plurality of fault tolerance processors **316** is configured to determine which one of the plurality of transaction processing servers **306** to designate as a primary instance. In particular, the plurality of fault tolerance processors **316** establish, upon the plurality of transaction processing servers **306** starting up, which of the plurality of transaction processing servers **306** is designated as the primary instance. Each of the others transaction processing servers **306** are then designated as a backup instance. Only one of the plurality of transaction processor servers **306** may be designated as the primary instance at the same time. One or more of the plurality of transaction processing servers **316** may be designated as backup instances. In some embodiments, different roles may be assigned by the plurality of fault tolerance processors **316** including a primary role or a backup role. Only one of the plurality of transaction processing servers **306** may be assigned with the primary role at the same time. More than one of the plurality of transaction processing servers **306** can be assigned the backup role.

(126) Each of the plurality of fault tolerance processors **316** is coupled with one transaction processing server of the plurality of transaction processing servers **306**. Each of the plurality of fault tolerance processors **306** is configured to determine whether a failure has occurred in the transaction processing server **306** designated as the primary instance and, subsequent thereto, cause one of the others of the plurality of transaction processing servers **306** to be designated as a new primary instance.

(127) In one embodiment, each of the plurality of tolerance processors **316a**, **316b**, and **316c** is configured to notify the other plurality of fault tolerance processors **316a**, **316b**, and **316c** if a failure in the transaction processing server designated as the primary instance **306a** occurred.

(128) In one embodiment, the failure in the primary instance **306a** may include a component failure in any one of the plurality of components **314a** of the primary instance **306a**. Each of the plurality of components **314a** is configured to notify the corresponding fault tolerance processor **316a** that a component failure occurred.

(129) In one embodiment, one or more of the plurality of components **314** of the transaction processing server **306** designated as the primary instance are each configured to communicate an electronic result message indicative of a result of an operation thereof to the corresponding component **314** of each of the other transaction processing servers **306** designated as the backup instance. Subsequent to the receipt of the electronic result message indicative of the result of the operation, the corresponding component **314** of each of the other transaction processing servers **306** designated as the backup instance compares a result of the operation thereof with the result indicated by the received message to determine whether one of the plurality of components **314** of the transaction processing server **306** designated as the primary instance or the corresponding component **314** of the other transaction processing servers **306** designated as the backup instance is operating correctly.

(130) In one embodiment, each of the plurality of fault tolerant processors **316** periodically

receives a communication from the others of the plurality of fault tolerant processors **316**. A failure to receive the communication from one of the plurality of fault tolerant processors **316** is deemed to indicate a failure of the transaction processing server **306** coupled therewith.

(131) In one embodiment, each of the plurality of components **314** of each of the plurality of transaction processing servers **306** are further coupled with the fault tolerant processor **316** coupled with that transaction processing server **306** such that the fault tolerant processor is able to detect a failure of any of the plurality of components coupled therewith.

(132) In one embodiment, the state of each of the plurality of components of each of the other transaction processing servers designated as the backup instance corresponds to the state of each of the plurality of components of the transaction processing server designated as the primary instance subsequent to the processing of an incoming request preceding the incoming request currently being processed by the transaction processing server designated as the primary instance.

(133) In one embodiment, each of the plurality of components **314** of each of the plurality of transaction processing servers **306** registers with a corresponding fault tolerance processor **316**, i.e., a fault tolerance group. Each of the plurality of fault tolerance processors **316** determines which one of the plurality of transaction processing servers **306** to designate as the primary instance by determining which one of the plurality of transaction processing servers **306** is the first to communicate an electronic message indicative of the current state in which all of the plurality of components **314** have been registered with the corresponding fault tolerance processor **316**.

(134) In one embodiment, as shown in FIG. 3, the system **300** includes a plurality of transaction processing servers/instances including a first, second, and third transaction processing servers **306a**, **306b**, and **306c** which are coupled with each other via the electronic communications network. The system **300** further includes a plurality of fault tolerance processors/groups including a first, second and third fault tolerance processor/FT **316a**, **316b**, and **316c** which are coupled/connected with each other via the electronic communications network. In one embodiment, the connection between the plurality of transaction processing servers and fault tolerance processors may be a bus, communications network or other electronic communications/data transfer medium, and the like. As used herein, bus refers to any communications system which transfers data between software and/or hardware components within a computer or between computers, and the like, and refers to all related hardware components, e.g., wire, optical fiber, and the like, and software, including communication protocols, and the like. The first, second, and third transaction processing server **306a**, **306b**, and **306c** may be coupled with the first, second, and third fault tolerance processor **316a**, **316b**, and **306c** respectively via the electronic communications network.

(135) The first, second, and third transaction processing servers **306a**, **306b**, and **306c** may each include a plurality of components **314a**, **314b**, and **314c** respectively, which are configured to register with the corresponding first, second, and third fault tolerance processor **316a**, **316b**, and **316c** which are configured to notify each other when the plurality of components **314a**, **314b**, and **314c** have registered successfully.

(136) Each of the plurality of components **314a**, **314b**, and **314c** that are registered with, i.e., that belong to the same transaction processing server **306a**, **306b**, or **306c** are coupled with each other via various interconnects. In one embodiment, each of the interconnects of the plurality of components **314a**, **314b**, and **314c** may comprise a bus or other inter-process/inter-component communications/data transfer mechanism such as a shared memory, and the like. These interconnects may be implemented with message queues which provide buffering between the components, allowing the components to operate independently/asynchronously to one another, e.g., at different processing rates.

(137) Each of the plurality of components **314a**, **314b**, and **314c** is characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an

operation which may alter the current state thereof.

(138) As further described below, each of the first, second, and third fault tolerance processors **316a**, **316b**, and **316c** periodically receives a communication from each other of the first, second and third fault tolerance processors **316a**, **316b**, and **316c**. A failure to receive the communication by the first, second, or third fault tolerance processor **316a**, **316b**, and **316c** from the others first, second, or third fault tolerance processors **316a**, **316b**, or **316c**, is deemed to indicate a failure of the corresponding transaction processing server **306a**, **306b**, or **306c**.

(139) In one embodiment, each of the first, second, and third plurality of components **314a**, **314b**, and **314c** are further coupled with the corresponding fault tolerant processor **316a**, **316b**, and **316c** such that the fault tolerant processor is able to detect a failure of any of the plurality of components **314**.

(140) In one embodiment, the first, second, and third fault tolerance processors **316a**, **316b**, and **316c** vote amongst each other to designate the first transaction processing server **306a** as the primary instance **306a** since it is the first one to report that the plurality of components **314a** have registered. The second and third transaction processing servers **306b** and **306c** are designated as backup instances **306b** and **306c**.

(141) In one embodiment, the primary instance **306a** is the only instance that receives an incoming request from a recipient external thereto, e.g., a market participant or a customer, and that transmits the generated electronic message indicative of the produced result to the recipient external thereto. The backup instances **306b** and **306c** are configured to receive state messages, e.g., state replication messages, from the primary instance **306a**, such as messages comprising the incoming requests received thereby. In particular, each of the plurality of components **314a** of the primary instance **306a** is configured to transmit result messages indicative of a result of an operation thereof to the corresponding components **314b** and **314c** of the backup instances **306b** and **306c**. Subsequent to the receipt of the result message, the corresponding components **314b** and **314c** of each of the backup instances **306b** and **306c** compares a result of the operation thereof with the result indicated by the result message to determine whether one of the plurality of components **314a** of the primary instance **306a** or one of the corresponding components **314b** and **314c** of the backup instances **306b** and **306c** is operating correctly. In particular, the component **314a** of the primary instance **306a** is configured to process an electronic message and to generate a primary result output. Each component **314b** and **314c** of each of the backup instances **306b** and **306c** is configured to receive the electronic message from the component **314a** of the primary instance **306a**, process the electronic message, and generate a backup result output. Each component **314b** and **314c** receives the output result message from its primary component counterpart **314a** and reconciles the backup result output with the primary result output of the primary component counterpart **314a** to determine correct operation.

(142) In one embodiment, the state of each of the plurality of components **314b** and **314c** of each of the backup instances **306b** and **306c** corresponds to the state of each of the plurality of components **314a** of the primary instance **306a** subsequent to the processing of an incoming request preceding the incoming request currently being processed by the primary instance **306a**.

(143) FIG. 5 depicts a detailed block diagram of the integrated component and fault tolerant system **300** according to one embodiment.

(144) In one embodiment, the system **300** may include a first, second, and third fault tolerance processors/FT **316a**, **316b**, **316c** coupled with the respective first, second, and third instances **306a**, **306b**, and **306c**.

(145) In one embodiment, the system **300** may include a plurality of instances/servers/groups **306a**, **306b**, and **306c** which are separate from each other. In one embodiment, each instance/server operates in its own physical hardware. In one embodiment, the first, second, and third fault tolerance processors **316a**, **316b**, and **316c** designate the first transaction processing server **306a** as a primary instance/server/group **306a**, the second transaction processing server **306b** as a backup

instance/server/group **306b**, and the third transaction processing server **306c** as a second backup instance/server/group **306c**.

(146) Each of the primary instance **306** and the backup instance **306b** may include a plurality of core components **314a** and **314b** respectively such as the plurality of core components **314a** and **314b** shown in FIG. 3. The second backup instance **306c** may also include a plurality of components **314c** (not shown in FIG. 5).

(147) As mentioned above, even though the embodiments discussed below with respect to FIGS. 5-20, may only refer to a single backup instance, it will be appreciated that more than one backup instance may be implemented such as the second backup instance **306c** which includes the plurality of components **314c** (not shown in FIG. 5).

(148) Synchronization of Primary Instance and Backup Instance

(149) In one embodiment, as shown in FIG. 5, a customer **318** may submit and transmit a transaction message, such as a trade request, via a request message flowing into the primary instance **306a** via the electronic communications network using, for example, the TM interface.

(150) The primary instance **306a** is configured to receive the request message, queue the request message, and route the request message to the transaction processor component **314b2** of the backup instance **306b** for backup processing and matching. The primary instance **306a** also routes the results of the processing and matching by its transaction processor component **314a2** to the backup instance **306b** for backup reconciliation. The result generator component **314a3** of the primary instance **306a** generates and queues result, e.g., market data, messages for publication and transmission to other customers/market participants and sends the result/market data messages to the result generator component **314b3** of the backup instance **306a**.

(151) In one embodiment, the plurality of components **314a** and **314b** of the primary instance **306a** and the backup instance **306b** may include a first, second, third, and fourth component. The first component may include a TR. The TR may be implemented by a TR processor such as the processor **202** described above with respect to FIG. 2. The second component may include a TP. The TP may be implemented by a TP processor such as the processor **202** described above with respect to FIG. 2. The third component may include an RG. The RG may be implemented by an RG processor such as the processor **202** described above with respect to FIG. 2. The fourth component may include a TL. The TL may include a database for storing the logged transaction and other data and may be implemented by a TL processor such as the processor **202** described above with respect to FIG. 2.

(152) As previously noted, communications between the plurality of components **314a** and **314b** of the primary instance **306a** and the backup instance **306b** may take place via an inter-instance communications bus/network which interconnects the instance **306a** and the backup instance **306b** and may be implemented, at least in part, by the network **220** shown in FIG. 2 and described elsewhere herein.

(153) In particular, corresponding first components **314a1** and **314b1** of the plurality of components **314** of the primary instance **306a** and the backup instance **306b** may include a TR **314a1** and **314b1** respectively. Corresponding second components **314a2** and **314b2** of the plurality of components **314a** of the primary instance **306a** and the backup instance **306b** may include a TP **314a2** and **314b2** respectively. Corresponding third components **314a3** and **314b3** of the plurality of components **314** of the primary instance **306a** and the backup instance **306b** may include a RG **314a3** and **314b3** respectively. Corresponding fourth components **314a4** and **314b4** of the plurality of components **314** of the primary instance **306a** and the backup instance **306b** may include a TL **314a4** and **314a4** respectively.

(154) FIG. 6 depicts a diagram of the inbound message flow to the primary instance **306a** shown in FIG. 5. As noted above, all communication between components of the same instance may be performed via IPC communication (shared memory). All communication between instances may be performed via the electronic communications network **220**.

(155) Once the negotiation of which instance is the primary instance and which one is the backup instance is complete, the first fault tolerance processor **316a** sends a feedback message to all components **314a**, including the TR **314a1**, the TP **314a2**, the RG **314a3**, and the TL **314a4**, indicating that each component **314a** is part of a primary instance and therefore, the messages may be transmitted to external recipients outside of the system. Similarly, the second fault tolerance processor **316b** sends a feedback message to all components **314b**, including the TR **314b1**, the TP **314b2**, the RG **314b3**, and the TL **314b4**, indicating that each component **314b** is part of a backup instance and, therefore, the output messages should not be transmitted to external recipients.

(156) As noted above, interconnects between components may be implemented with message queues, e.g., bus(es) as shown in FIGS. 5-9 which provide buffering between the components. In one embodiment, the customer **318** may submit and transmit an electronic request message (Transaction Input message **330**) including data indicative of a trade, via the electronic communications network **220**, flowing into the primary instance **306a** via the TR **314a1**. The TR **314a1** is configured to receive the electronic request message, send confirmation to the customer **318** that the electronic request message has been received, queue the electronic request message, and route, via shared communication, the electronic message to the TP **314a2** for processing. The TP **314a2** is configured to queue a copy of the unprocessed request message (TP input replication/replicated message **330**) and to route the copy of the unprocessed request message to the backup instance **306b** via the electronic communications network. The TP **314a2** is configured to queue the generated result of the processing and to route the result to the TR **314a1** and to the RG **314a3**. The RG **314a3** is configured to generate and queue output market data messages for publication and transmission to other customers/market participants outside of the system. The TR **314a1** is configured to process the result message indicative of the produced result and to queue the result message. The TR **314a1** and the RG **314a3** sends the result message and the RG generated messages to the TL **314a4** for storage and consumption by the regulatory component.

(157) In one embodiment, the queues/buses provide the ability to determine the point of failure. For example, if there is an input message that causes a failure in the TP **314a2**, the system **300** is able to isolate, recover from the failure, and avoid a failure in the backup TP **314b2** by using the queues to identify and avoid processing that same input message. In particular, in one implementation, when only a single component is determined to have failed, the input message determined to have been in process by that failed component at the time of failure, as determined based on the those input messages for which a result was produced by the failed component prior to failure, is identified to an operator of the system and recovery by the backup instance is paused until the operator determines that the in-process message was not the cause of the failure.

(158) In one embodiment, the system **300** is configured to ensure that the TP **314a2** in the primary instance **306a** fully processes the request message successfully before the TP **314b2** in the backup instance **306b** starts processing request message. Therefore, the system **300** provides protection against faulty/corrupted messages which may cause failures in the different components. Further, only successfully processed request messages are stored into the TL database **314a4**.

(159) FIG. 7 depicts a block diagram of the input message flow into the backup instance **306b** of the system **300** shown in FIG. 3. In general, the primary and backup instances **306a** and **306b** are identical and include the same components arranged and interconnected in the same manner, e.g., the TR **314b1**, the TP **314b2**, the RG **314b3**, and the TL **314b4** operate in the same manner as their primary instance counterparts e.g., the TR **314a1**, the TP **314a2**, the RG **314a3**, and the TL **314a4**. That is, a given instance is capable of operating as either primary or a backup. However, as discussed above, there are a few differences in the way an instance operates depending on whether it is configured as a primary or as a backup. For example, the TR **314b1** receives the incoming request message from the TP **314a2** instead of receiving the incoming request message from the customer **318**. Further, as already indicated, each of the components of the backup instance **306b** waits until its counterpart in the primary instance **306a** has successfully performed the

corresponding operation before starting the processing in the backup instance **306b**. After all components **314b** of the backup instance **306b** have processed the copy of the electronic message, then the state of the backup instance **306b** should be one step behind the primary instance **306a**. Further, the TR **314b1** and the RG **314b3** do not communicate any messages to external recipients. (160) FIG. **8** depicts a block diagram of the outbound message flow from the primary instance **306b** of the system **300** according to one embodiment. FIG. **9** depicts a block diagram of the outbound message flow from the backup instance of a fault tolerant system according to one embodiment.

(161) In one embodiment, referring to FIGS. **5-9**, a first component of the plurality of components **314** of the transaction processing server designated as the primary instance **306a** and each of the transaction processing servers designated as the backup instance **306b** comprises a TR **314a1** and **314b1** respectively for receiving the electronic request message **330** via the electronic communications network, the electronic request message **330** including data indicative of an order to trade. As mentioned above, the TR **314b1** receives the incoming request message **330** from the TP **314a2** instead of receiving the incoming request message from the customer **318**. A second component of the plurality of components of the transaction processing server designated as the primary instance **306a** and each of the plurality of transaction processing servers designated as the backup instance **306b** comprise a TP **314a2** and **314b2** respectively for receiving and processing the incoming request **330** to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request **330** and the previously received request.

(162) The TR **314a1** of the transaction processing server designated as the primary instance **306a** is configured to route, via a shared memory, the received request message **330** to the TP **314a2** of the primary instance **306a**. In particular, the received electronic message **330** is queued and then routed to the TP **314a2**.

(163) The TP **314a2** of the transaction processing server designated as the primary instance **306a** is configured to: receive, via the shared memory, the incoming request message **330**, forward, upon receipt of the incoming request message **330**, via the electronic communications network, the incoming request message **330** to the TR **314b1** of each of the plurality of transaction processing servers designated as the backup instance **306b**, process the received incoming request **330** to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request **330** and the previously received request and produce a result **332** based thereof; and forward, via the electronic communications network, the electronic result message **332** to the TP **314b2** of each of the plurality of transaction processing servers designated as the backup instance **306b**.

(164) The TR **314b1** of each of the transaction processing servers designated as the backup instance **306b** is configured to: receive, via the electronic communications network, the incoming request **330** from the TP **314a2** of the transaction processing server designated as the primary instance **306a**; and forward, via a shared memory, the incoming request **330** to the TP **314b2** of the corresponding transaction processing server designated as the backup instance **306b** for backup processing. As shown in FIG. **9**, the TR **314b1** compares TR primary **334a** and backup **334b** outputs/result messages to determine whether the TR **314a1** is operating correctly.

(165) The TP **314b2** of each of the transaction processing servers designated as the backup instance **306b** is configured to: receive, via a shared memory, the incoming request **330** from the TR **314b1** of the corresponding transaction processing server designated as the backup instance **306b**; receive, via the electronic communications network, the electronic result message indicative of a primary result from the TP of the transaction processing server designated as the primary instance **306a**; perform backup processing of the incoming request **330**, in response to the receipt of the incoming request **330** and the electronic result message (TP Output from primary **332a**), wherein the backup processing includes generation of an electronic backup result message (TP output **332b**) indicative

of a backup result; and compare, as shown in FIG. 9, the primary result **332a** with the backup result **332b** to determine whether the TP **314a2** of the transaction processing server designated as the primary instance **306a** is operating correctly. In other words, the TP **314b2** reconciles the TP outputs **332a** and **332b** produced respectively by the TP **314a2** of the primary instance and by the TP **314b2** of the backup instance.

(166) In one embodiment, a third component of the plurality of components **314a** of the transaction processing server designated as the primary instance **306a** and each of the other transaction processing servers designated as the backup instance **306b** comprise a RG **314a3** and **314b3** for generating and publishing (only by the RG **314a3**) a market data message. The RG **314a3** of the transaction processing server designated as the primary instance **306a** is configured to: receive, via the shared communication, from the TP **314a2** of the transaction server designated as the primary instance **306a**, the electronic result message; generate a market data message (RG Output from primary **336a**) based on the electronic result message; forward, via the electronic communications network, the market data message **336a** to the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**; and distribute, via the electronic communications network, the market data message to the plurality of market participants.

(167) The RG **314b3** of each of the other transaction processing servers designated as the backup instance is configured to: receive, via the communications network, from the RG **314a3** of the transaction processing server designated as the primary instance **306a**, the market data message **336a**; receive, via a shared memory, from the TP **314b2** of the transaction processing server designated as a backup instance **306b**, the backup TP output result **332b**; generate a backup market data message (RG Output **336b**), in response to receipt of the market data message, based on the backup result; and compare, as shown in FIG. 9, the backup market data message **336b** with the primary market data message **336a** to determine whether the RG **314a3** is operating correctly. In other words, the RG **314b3** reconciles the outputs RG output **336a** and **336b** produced by the RG **314a3** of the primary instance **306** and by RG **314b3** of the backup instance.

(168) Startup Sequence of Primary Instance

(169) Both the primary and backup components include identical computer executable instructions, but the FT role differentiates how the component operates while assigned the primary role or the backup role.

(170) FIGS. **10A** and **10B** depict a sequence diagram of a startup sequence of a transaction processing server **306a**, assigned with a primary instance role, of the system **300** shown in FIG. 3.

(171) At startup, the FT **316a** is configured to startup before the plurality of components/applications **314a** may start. The FT **316a** is configured to wait to receive a request from each component to join and register. Once the FT **316a** has started, the core components **314a** including the match engine **314a2**, the TR **314a1**, the RG **314a2**, and the TL **314a4** components may startup in any order. At startup, the core components **314a** may initialize, starting with loading static configurations independently from each other and initiating shared memory communication IPC and network communication endpoints. Once each component **314a** starts, each component is configured to communicate with the corresponding FT **316a** via the electronic communications network, e.g., via a TCP connection. If any single component **314a** fails at this stage, the failed component is restarted. Each component **314a** sends a request to join the group and register with the FT **316a**. When the FT **316a** determines that the four components have joined and registered, then the group is finalized. The FT **316a** communicates with the other FTs **316b** and **316c** and starts voting to assign the instance/group with a fault tolerance role of either a primary or a backup role. The FT **316a** determines that there is no other instance with a primary role assigned and notifies each registered component **314a** residing on the same group or server that they are assigned the primary role. Once each registered component **314a** knows that they have been assigned with the primary role, each component **314a** finalizes their startup process as a primary component. If any single component **314a** is stuck not receiving the primary role, the failed component is restarted.

Depending on the role of each component **314a**, the startup process is different. Each component **314a** starts up independently. However, the components **314a** are synchronized via a message status sent by the TR **314a1** via shared memory communication to notify the other components such as the TP **314a2**, the RG **314a3**, and TL **314a4** that the TR **314a1** has successfully started with the primary role. At this stage, this status message is the only message type that flows through IPC transport (before the FT-FINALIZE-PRIMARY event notification to the FT). Once each component receives the status message that the TR has initialized as primary, then each component may notify the FT **316a** that they are ready to start as a primary component (FT-FINALIZE-PRIMARY). The FT **316a** then transmits a PRIMARY-READY event message when it has received the FT-FINALIZE-PRIMARY event message from all four components.

(172) In one embodiment, the system **300** may support a configurable timeout for the PRIMARY-READY event call back notification from the time FT-FINALIZE-PRIMARY event is sent to the FT **316a**. The corresponding component **314a** may shut down if no response is received within this time out window.

(173) In one embodiment, the FT-PRIMARY-READY event call back message from the FT **316a** is not transmitted in a sequential order of client connections. Therefore, the components may handle messages being received on IPC transports until the FT-PRIMARY-READY call back is received from the FT **316a**. In this case, the component **314a** is configured to simply queue the messages received on the IPC transport and only process them after FT-PRIMARY-READY call back notification is received.

(174) Once all of the components **314a** are ready to start as primary components, they proceed together, and they can start processing the messages for customers with the primary role.

(175) Startup Sequence of Backup Instance

(176) FIGS. **11A-11D** depict a sequence diagram of a startup sequence of a transaction processing server **306b**, assigned with a backup instance role, of a system **300** including a plurality of startup stages according to one embodiment. In particular, FIG. **11A** depicts stages **1-2**, FIG. **11B** depicts stage **3**, FIGS. **11C(1)** and **11C(2)** depict stages **4-7**, and FIG. **11D** depicts stage **8**.

(177) As shown in FIG. **11A**, stage **1** includes the same sequence as the startup sequence of the primary instance **306a** discussed above. Stage **1** of the startup sequence includes static configurations and initialization. In particular, the FT **316b** is configured to startup before all the plurality of components/applications **314b** start. The FT **316b** is configured to wait to receive a request from each component **314b** to join and register. Once the FT **316b** has started, the core components **314b** including the match engine **314b2**, the TR **314b1**, the RG **314b3**, and the TL **314b4** components may startup in any order. At startup, the core components **314b** may initialize, starting with loading static configurations independently from each other and initiating shared memory communication IPC and network communication endpoints. Once each component **314b** starts, each component is configured to communicate with the corresponding FT **316b** via the electronic communication network, e.g., via a TCP connection. If any single component **314b** fails at this stage, the failed component is restarted. Each component **314b** sends a request to join the group and register with the FT **316b**. When the FT **316b** determines that the four components have joined and registered, then the group is finalized.

(178) Stage **2** includes finding the fault tolerance role. In particular, the FT **316b** communicates with the other FTs **316b** and **316c** and starts voting to assign the instance/group with a fault tolerance role of either a primary or a backup role. The FT **316b** determines that there is an existing instance with a primary role assigned and notifies each registered component **314b** residing on the same server that they are assigned the backup role.

(179) As shown in FIG. **11B**, in stage **3**, once each registered component **314b** knows that they have been assigned with the backup role, each component **314b** finalizes their startup process as a backup component and initialize message bus listeners via the electronic communications network to receive messages from the primary instance **306a**. In particular, the TR **314b1** of the backup

instance **306b** starts listening on the TP input bus for the inputs from the primary instance **306a**. The TP **314b2** starts listening on the TP output bus for the outputs of the primary instance **306a**. The RG **314b3** of the backup instance **306b** starts listening on the RG output bus for the outputs of the primary instance **306a**.

(180) As shown in FIGS. **11C(1)** and **11C(2)**, in stage **4**, the FT **316b** confirms with the different components **314b** that the components are ready to send the state requests via the call back messages (FT-FINALIZE-STATE-REQ) and (FT-STATE-REQ-READY).

(181) In stage **5**, the TR **314b1** of the backup instance **306b** sends a request for the for the current state of the primary instance **306a** to its counterpart (the TR **314a1**) of the primary instance **306a**. Some of the components **314b** (the TR **314b1**, the match engine **314b2**, and the RG **314b3**) are configured to synchronize with their counterpart in the primary instance **306a**. In particular, these components are configured to wait to receive the current state of their counterpart in the primary instance **306a**. In one embodiment, the TL **314b4** of the backup instance **306b** does not receive the state from its primary counterpart. In an alternative embodiment, the TL **314b4** of the backup instance **306b** receives state information from its primary counterpart so as be synchronized therewith as described herein.

(182) In Stage **6**, once the state of each component from the primary instance **306a** is received, each component TR **314b1**, match engine **314b2**, and RG **314b3** of the backup instance **306b** finish initialization based on the received state of the corresponding component from the primary instance **306a**.

(183) In Stage **7**, the components **314b** confirm with the FT **316b** of the backup instance **306b** that they are ready to receive an incoming transaction message, e.g., incoming trade order, processing, i.e., incoming request **330** message for processing. In particular, each component **314b** notifies the FT **316b** that they are ready to start as a component with the backup role by sending a “FT-FINALIZE-BACKUP” message. The FT **316b** then transmits a “FT-BACKUP-READY” when it has received the “FT-FINALIZE-BACKUP” event message from all four components.

(184) As shown in FIG. **11D**, in stage **8**, the components **314b** start processing messages in the backup role as described above.

(185) FIG. **13** illustrates an example flowchart **1300** of the operation of the system **300** shown in FIGS. **3** and **5**. In one embodiment, the operation of an integrated component system **300** includes: implementing, by a processor of each of a plurality of transaction processing servers **306** coupled with each other via an electronic communications network, a plurality of components **314** coupled with each other, each of the plurality of components characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof, wherein one of the plurality of transaction processing servers **306** is designated as a primary instance **306a** and each of the other transaction processing servers **306** is designated as a backup instance **306b**, wherein each of the plurality of transaction processing servers **306** is operative to receive an incoming request **330** for a transaction, process the received incoming request **330** to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request **330** and previously received request and generate a result indicative thereof (**1302**); at least commencing processing, by the transaction processing server **306a** designated as the primary instance, of the incoming request **330** prior to each of the other transaction processing servers designated as the backup instance **306b** (**1304**); determining, by each of a plurality of fault tolerance processors **316** coupled with one transaction processing server **306** of the plurality of transaction processing servers, whether a failure has occurred in the transaction processing server designated as the primary instance **306a** (**1306**); and causing, by each of the plurality of fault tolerance processors **316**, subsequent to determining that a failure occurred, one of the others of the plurality of transaction processing servers **316** to be designated as a new primary instance (**1308**).

(186) In one embodiment, determining, whether a failure has occurred includes communicating, by each of the one or more of the plurality of components **314a** of the transaction processing server designated as the primary instance **306a**, an electronic message indicative of a result of a primary component operation thereof to the corresponding backup component of each of the transaction processing servers designated as the backup instance **306b**; and comparing the result of the primary component operation with the result of the backup component operation.

(187) In one embodiment, the plurality of components **314** are coupled with each other via a shared memory architecture.

(188) In one embodiment, the operation of the integrated component system **300** further includes receiving, by each of the other transaction processing servers designated as the backup instance **306b**, the incoming request **330** from the transaction processing servers designated as the primary instance **306a**.

(189) In one embodiment, the state of each of the plurality of components **314b** of each of the other transaction processing servers designated as the backup instance **306b** corresponds to the state of each of the plurality of components **314a** of the transaction processing server designated as the primary instance **306b** subsequent to the processing of an incoming request preceding the incoming request currently being processed by the transaction processing server designated as the primary instance **306a**.

(190) In one embodiment, the operation of the integrated component system **300** further includes transmitting, only by the transaction processing server designated as the primary instance **306a**, a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and previously received request to a recipient via the electronic communications network.

(191) In one embodiment, the operation of the integrated component system **300** further includes establishing, by the plurality of fault tolerance processors **316**, upon the plurality of transaction processing servers **306** starting up, which of the plurality of transaction processing servers **306** is designated as the primary instance.

(192) In one embodiment, one or more of the plurality of transaction processing servers **306** is located in a geographic region different from a location of one or more others of the plurality of transaction processing servers **306**.

(193) In one embodiment, each of the plurality of components **314** of each of the plurality of transaction processing servers **306** are further coupled with the fault tolerance processor **316** coupled with that transaction processing server **306** such that the fault tolerance processor is able to detect a failure of any of the plurality of components **314** coupled therewith.

(194) In one embodiment, each of the plurality of fault tolerance processors **316** is configured to notify the other plurality of fault tolerance processors **316** if a failure of the transaction processing server designated as the primary instance has occurred.

(195) In one embodiment, the failure in the primary instance **306a** includes a component failure of any one of the plurality of components of the primary instance **314a**.

(196) In one embodiment, the operation of the integrated component system **300** further includes periodically receiving, by each of the plurality of fault tolerant processors **316**, a communication from the others of the plurality of fault tolerant processors **316**, wherein failing to receive, by each of the plurality of fault tolerant processors **316**, the communication from one of the plurality of fault tolerant processors **316** is deemed to indicate a failure of the transaction processing server **306** coupled therewith.

(197) In one embodiment, a first component of the plurality of components comprises a TR for receiving an incoming request **330** for a transaction via the electronic communications network. A second component of the plurality of components comprises a TP for receiving and processing the incoming request **330** to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request **330** and the previously received request. A third component of the plurality of components comprises an RG for generating and publishing a

market data message.

(198) In one embodiment, the operation of the integrated component system **300** further includes routing, by the TR **314a1** of the transaction processing server designated as the primary instance **306a**, via a shared memory, the received incoming request **330** to the TP **314a2** of the transaction processing server designated as the primary instance **306a**.

(199) In one embodiment, the operation of the integrated component system **300** further includes receiving, by the TP **314a2** of the transaction processing server designated as the primary instance **306a**, via a shared memory, the incoming request **330**; forwarding, by the TP **314a2** of the transaction processing server designated as the primary instance **306a**, via the electronic communications network **220**, the received incoming request to the TR **314a3** of each of the other transaction processing servers designated as the backup instance **306b**; processing, by the TP **314a2** of the transaction processing server designated as the primary instance **306a**, the received incoming request **330** to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and the previously received request and produce a result based thereof; generating, by the TP **314a2** of the transaction processing server designated as the primary instance **306a**, an electronic result message indicative of a produced result; and forwarding, by the TP of the transaction processing server designated as the primary instance **306a**, via the electronic communications network, the electronic result message to the TP **314b2** of each of the other transaction processing servers designated as the backup instance **306b**.

(200) In one embodiment, the operation of the integrated component system **300** further includes receiving, by the TR **314b1** of each of the other transaction processing servers designated as the backup instance **306b**, via the electronic communications network **220**, the incoming request **330** from the TP **314a2** of the transaction processing server designated as the primary instance **306a**; and forwarding, by the TR **314b1** of each of the other transaction processing servers designated as the backup instance **306b**, via a shared memory, the incoming request **330** to the TP **314b2** of the corresponding transaction processing server designated as the backup instance **306b** for backup processing.

(201) In one embodiment, the operation of the integrated component system **300** further includes: receiving, by the RG **314a3** of the transaction processing server designated as the primary instance **306a**, via a shared memory architecture, from the TP **314a2** of the transaction server designated as the primary instance **306a**, the electronic result message; generating, by the RG **314a3** of the transaction processing server designated as the primary instance **306a**, a market data message based on the electronic result message; forwarding, by the RG **314a3** of the transaction processing server designated as the primary instance **306a**, via the electronic communications network, the market data message to the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**; and distributing, by the RG **314a3** of the transaction processing server designated as the primary instance **306a**, via the electronic communications network **220**, the market data message to a plurality of market participants.

(202) In one embodiment, the operation of the integrated component system **300** further includes receiving, by the TP **314b2** of each of the other transaction processing servers designated as the backup instance **306b**, via a shared memory, the incoming request **330** from the TR **314b1** of the corresponding transaction processing server designated as the backup instance **306b**; receiving, by the TP **314b2** of each of the other transaction processing servers designated as the backup instance **306b**, via the electronic communications network, the electronic result message indicative of a primary result from the TP **314a2** of the transaction processing server designated as the primary instance **306a**; performing, by the TP **314b2** of each of the other transaction processing servers designated as the backup instance **306b**, backup processing of the incoming request **330**, in response to the receipt of the incoming request **330** and the electronic result message, wherein the backup processing includes generation of an electronic backup result message indicative of a backup result; and comparing, by the TP **314b2** of each of the transaction processing servers

designated as the backup instance **306b**, the primary result with the backup result to determine whether the TP **314a2** of the transaction processing server designated as the primary instance **306a** is operating correctly.

(203) In one embodiment, the operation of the integrated component system **300** further includes receiving, by the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**, via the communications network, from the RG **314a3** of the transaction processing server designated as a primary instance **306a**, the market data message; receiving, by the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**, via a shared memory, from the TP **314b2** of the corresponding transaction processing server designated as a backup instance **306b**, the backup result; generating, by the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**, a backup market data message, in response to receipt of the market data message, based on the backup result; and comparing, by the RG **314b3** of each of the other transaction processing servers designated as the backup instance **306b**, the backup market data message with the market data message to determine whether the RG of the transaction processing server designated as the primary instance **306a** is operating correctly

(204) Fault Tolerant System

(205) As noted above, the system **300** provides fault tolerance features to be able to handle fault tolerance in an integrated component system environment.

(206) FIG. **12** depicts a sequence diagram of a failover sequence of a backup instance **306b** in a fault tolerant system **300** as shown in FIG. **3**.

(207) In one embodiment, the system **300** includes a plurality of transaction processing servers **306** coupled with each other via an electronic communications network. Each of the plurality of transaction processing servers **306** includes a processor which implements a plurality of components **314** coupled with each other. Each of the plurality of components **314** is characterized by a current state and is configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof. One of the plurality of transaction processing servers **306** is designated as a primary instance **306a** and each of the other transaction processing servers **306** is designated as a backup instance **306b**. Each of the plurality of transaction processing servers **306** is operative to receive an incoming request for a transaction, process the received request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and previously received request and generate a result indicative thereof. The transaction processing server designated as the primary instance **306a** at least commences processing of the incoming request prior to each of the other transaction processing instances designated as the backup instance. The one or more of the plurality of components **514a** of the transaction processing server designated as the primary instance **306a** are each configured to communicate an electronic result message indicative of a result of an operation thereof to the corresponding component **314b** of each of the other transaction processing servers designated as the backup instance **306b**.

(208) In one embodiment, only the transaction processing server designated as the primary instance **306** is configured to transmit a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and previously received request to a recipient external to the system.

(209) The system **300** further includes a plurality of fault tolerance processors **316a**, **316b**, and **316c**. Each of the plurality of fault tolerance processors **316a**, **316b**, and **316c** is coupled with one transaction processing server **306a**, **306b**, and **306c** of the plurality of transaction processing servers **306** respectively. Each the plurality of fault tolerance processors **316a**, **316b**, and **316c** is configured to determine that a failure has occurred in the transaction processing server designated as the primary instance **306a** and, subsequent thereto, cause one of the others of the plurality of

transaction processing servers **306b** to be designated as a new primary instance. The transaction processing server which is designated as the new primary instance **306b** is configured to synchronize the current state of each of the plurality of components **314b** thereof based on the received electronic messages indicative of a result of the operation the corresponding component **314a** of the primary instance **306a**, and process any incoming requests not processed by the primary instance **306a** before commencing processing of new incoming requests.

(210) In one embodiment, the state of each of the plurality of components **314b** of each of the other transaction processing servers designated as the backup instance **306b** corresponds to the state of each of the plurality of components **314a** of the transaction processing server designated as the primary instance **306a** subsequent to the processing of an electronic transaction message preceding the electronic transaction message currently being processed by the transaction processing server designated as the primary instance **306a**.

(211) In one embodiment, the plurality of fault tolerance processors **316**, establish, upon the plurality of transaction processing servers starting up, which of the plurality of transaction processing servers is designated as the primary instance **306a**.

(212) In one embodiment, each of the plurality of components **314** of each of the plurality of transaction processing servers **306** are further coupled with the fault tolerance processor **316** coupled with that transaction processing server such that the fault tolerance processor is able to detect a failure of any of the plurality of components.

(213) In one embodiment, only a failure in the transaction processing server **306a** or the fault tolerance processor **316a** of the transaction processing server designated as the primary instance causes one of the others of the plurality of transaction processing servers **306b** to be designated as a new primary instance **306a**. In other words, only a failure in the primary instance **306a** or the fault processor **306b** causes a failover process to occur. If a transaction processing server designated as the backup instance **306b**, the corresponding fault tolerance processor **316b**, or any of the corresponding components **314b** fail, the transaction processing server designated as the primary instance **306a** continues to operate without that backup instance **306b** and the failed backup instance **306b** may be brought down. In one embodiment, an operator of the system may be notified to recognize that another backup instance **306b** needs to be started, particularly if the failed backup instance **306b** was the only backup instance **306b**. In an alternative embodiment, where all available backup instances **306b** have failed, the fault tolerance processor **316a** of the primary instance **306a** may shut down the primary instance **306a**, e.g., performing a graceful shutdown where the operator or traders are notified and all in process transactions are completed, so as to minimize the possibility of lost transactions should the primary instance **306a** fail with no backup instance **306b** synchronized therewith and ready to take over.

(214) Failures in the primary instance **306a** may occur due to a failure on one of the components in the primary instance **306a**, a failure in the fault tolerance processor **316a**, a failure in the network, e.g., preventing communications with the primary instance **306a**, a failure in communication with the transaction processing server **306a** designated as the primary instance, a failure due to a message that cannot be processed by the transaction processing server **306a** designated as the primary instance or which causes a processing fault, e.g., a faulty message, a catastrophic failure in the transaction processing server **306a** designated as the primary instance, e.g., the hardware server box fails due to a power outage or other catastrophic event, and the like.

(215) In one embodiment, each of the plurality of fault tolerance processors/FT **316a**, **316b**, and **316c** is configured to determine whether a failure occurred in the transaction processing server designated as the primary instance **306a**. In other words, a failure that occurs in the primary instance **306a** may be detected by any of the plurality of fault tolerance processors **316a**, **316b**, and **316c**.

(216) In one embodiment, each of the plurality of tolerance processors **316a**, **316b**, and **316c** is configured to notify the other plurality of fault tolerance processors **316a**, **316b**, and **316c** if a

failure in the transaction processing server designated as the primary instance **306a** occurred.

(217) In one embodiment, the failure in the primary instance **306a** may include a component failure in any one of the plurality of components **314a** of the primary instance **306a**.

(218) In one embodiment, each of the plurality of components **314a** is configured to, when capable, notify the corresponding fault tolerance processor **316a** that it has failed. For example, where a component **314a** is unable to complete an external operation, such as database commit, or otherwise detects an internal fault, it may notify the corresponding fault tolerance processor **316a** that it is unable to continue to operate properly. In one embodiment, the failed component intentionally, or unintentionally as a byproduct of the failure, may close some or all shared memory and network communications and proceed to shut down, or await being shut down, which may then be detected, e.g., via being non-responsive or via the detection of the closure of a communications channel, by the corresponding fault tolerance processor **316a** as an indication of failure. As described herein, active notification by a component of its failure may be combined with the passive failure detection described herein, e.g., via detection of a lost, dropped or closed connection, failure to receive a heartbeat communication within a threshold period of time, and/or failure to respond to a communication necessitating a reply within a threshold time period, etc.

(219) In one embodiment, each of the plurality of fault tolerant processors **316a**, **316b**, and **316c** periodically receives a communication from the others of the plurality of fault tolerant processors. A failure to receive the communication from one of the plurality of fault tolerant processors **316a**, **316b**, and **316c** is deemed to indicate a failure of the transaction processing server **306** coupled therewith. For example, a failure to receive the communication from the fault tolerant processor **316a** which is coupled with the primary instance **306c** is deemed to indicate a failure in the primary instance **306a**.

(220) In one embodiment, the communication may include a heartbeat message. If a heartbeat message is not received by a threshold time, e.g., 10 seconds, a failure is deemed to have occurred.

(221) In one embodiment, upon receipt of the electronic result message indicative of the result of the operation of the one or more of the plurality of components **314a** of the primary instance **306a**, the corresponding component **314b** of each of the other transaction processing servers designated as the backup instance **306b** compares a result of the operation thereof with the result indicated by the received electronic result message to determine whether one of the plurality of components of the transaction processing server designated as the primary instance or the corresponding component of each of the other transaction processing servers designated as the backup instance is operating correctly.

(222) In one embodiment, when a fault tolerance processor of the plurality of fault tolerance processors **316a**, **316b**, or **316c** determines that a failure occurred in the primary instance **306a**, the plurality of fault tolerance processors **316a**, **316b**, or **316c** are configured to designate one of the others of the plurality of transaction processing servers as a new primary instance. In particular, once a failure occurs in the primary instance **306a**, the fault tolerance processors **316a**, **316b**, and **316c** are configured to vote and agree on a new primary instance.

(223) In one embodiment, as shown in FIGS. **12A(1)**, **12A(2)**, **12A(3)**, and **12B**, the backup instance **316b** is designated by the plurality of fault tolerance processors **316a**, **316b**, and **316c** as the new primary instance. The FT **316b** may notify each registered component **314b** residing on the same server that they will become part of the new primary instance and may transmit a "PRIMARY" message.

(224) In one embodiment, all instances including the primary instance **306a** and the backup instance **306b** may include a queue/bus configured to store previously received messages. The backup instance **306b** that is designated as the new primary instance is configured to retrieve the stored request messages from the queue and process all of the stored request messages, generate an electronic result message indicative of a backup output result based on the processing, and synchronize all of the plurality of components **314b** before becoming the new primary instance.

(225) In one embodiment, the electronic message indicative of the result of the operation of the transaction processing server designated as the primary instance **306a** includes a primary result output. The electronic result message indicative of the result of the operation of each of the plurality of transaction processing servers designated as the backup instance includes a backup result output.

(226) In one embodiment, at least a subset of the plurality of components **314b** of the transaction processing server **306b** designated as the new primary instance are configured to complete a reconciliation of the backup result output with the primary result output.

(227) In one embodiment, at least a subset of the plurality of components **314b** of the transaction processing server **306b** designated as the new primary instance are configured to identify a backup result output that does not have a respective primary instance result output and publish an unmatched result output via the shared memory communication. This unmatched result output may be communicated to the fault tolerance processor of the transaction processing server **306b** designated as the new primary instance and/or to a system operator to determine whether the unmatched output is based on the failure from which recovery is occurring and to determine whether any reprocessing is needed.

(228) In one embodiment, when the fault tolerance processor **316a** of the plurality of fault tolerance processors **316** coupled with the primary instance **316a** determines that the failure has occurred in the primary instance, each of the plurality of fault tolerance processors **316** is configured to cause the primary instance to at least de-couple from the others of the plurality of transaction processing servers. In one embodiment, each of the plurality of fault tolerance processors **316** is configured to cause the primary instance to de-couple from the electronic communications network **220**.

(229) In one embodiment, the plurality of components **314b** may include a TR **314b1**, a TP **314b2**, an RG **314b3**, and a TL **314b4**. Subsequent to receiving the notification to become a primary component, the plurality of components **314b** may perform pre-requisite tasks in preparation to take over as a primary component. In particular, before the transaction processing server **306b** designated as the new primary instance may takeover, the plurality of components **314b** need to be synchronized. In particular, as described below, coordination between the plurality of components **314b** may be accomplished via status messages via shared memory communication.

(230) Therefore, a synchronizing process described below coordinates all of the plurality of components **314b** of the transaction processing server **306b** designated as the new primary to process any electronic messages received before and after the failure occurred.

(231) In one example, the TR **314a1** may be processing message **9**, while the TP **314a2**, is processing message **8**, the RG **314a3** is processing message **6**, and the TL **314a4** is processing message **4**. If the TR **314a1** fails when processing message **9**, then it notifies the other components so they know they will not receive message **9**. The TR **314a1** also notifies the transaction processing server **306a** that a failure occurred at message **9**. The transaction processing server **306a** designated as primary instance continues to queue messages until all components on the transaction processing server **306b** designated as the new primary instance have caught up in processing messages **1-8** before it can become the primary instance and creates a save state. This ensures that all messages are being processed sequentially and all messages have been processed before failing over to the transaction processing server **306b** designated as the new primary instance.

(232) In particular, at this stage, the TR **314b1** leads the coordination of the plurality of components **314b**. The TR **314b1** may continue to listen for copies of electronic messages (TP input **330**) and flush all queued inputs including copies of electronic messages previously received from the failed primary instance **306a**. The TP **314b2**, the RG **314b3**, and the TL **314b4** may continue to process any queued inputs as a backup instance until a takeover message is received, e.g., a “NCSS takeover” from the TR **314b1**.

(233) Once all queued inputs from the failed primary instance **306a** have been flushed, the TR **314b1** may notify the TP **314b2** and may transmit a takeover message, e.g., a “NCSS takeover”.

(234) When the takeover message is received, some of the plurality of components **314b** of the transaction processing server **306b** designated as the new primary instance may be configured to complete reconciliation of backup result outputs with primary result outputs, identify backup result outputs that do not have respective primary instance result outputs, and publish unmatched result outputs via the shared memory. The reconciliation of backup result outputs enables the system **300** to confirm that the transaction server designated as the new primary instance backup is in a healthy state, that it has caught up with the primary instance up until the point of failure, and that it is safe for the transaction processing server **306b** to take over as the new primary instance. Where the reconciliation fails, the backup instance **306b** may be determined to have failed and the recovery process then begins again with another backup instance. Alternatively, it may be determined that recovery is not possible.

(235) In particular, at this stage, the TP **314b2** leads the coordination of the plurality of components **314b**. The TP **314b2** is configured to complete reconciliation of corresponding backup result outputs with primary result outputs, identify backup result outputs that do not have respective primary instance outputs, and publish unmatched outputs via the shared memory communication to the other plurality of components **314b**, e.g., to identify those transactions which must be reprocessed thereby. Further, the TP **314b2** publishes the unmatched outputs via the electronic communication network **220**. The TP **314b2** is configured to transmit, via the shared memory, the takeover message to the other plurality of components **314b** including the TR **314b1**, the RG **314b3**, and the TL **314b4**. When the takeover message is received, the TR **314b1** and the RG **314b3** are also configured to complete the same tasks including completing reconciliation of corresponding backup outputs with primary outputs, identify backup outputs that do not have respective primary instance outputs, and publish unmatched result outputs via the shared memory. When the takeover message is received, the TL **314b4** does not perform any reconciliation. Instead, the TL captures **314b4** messages via the shared memory communication and acknowledges the takeover.

(236) Each of the plurality of components **314b** of the transaction processing server **306b** designated as the new primary instance may notify the fault processor **316a** to finalize the takeover as the new primary instance via a finalize takeover message, e.g., “FINALIZE-TAKEOVER.”.

(237) When the fault processor **316b** receives the finalize takeover message from each of the plurality of components **314b**, the fault processor **316b** notifies each of the plurality of components **314b** that each of the plurality of components **314b** is ready to take over as a new primary instance by transmitting a takeover ready message, e.g., “TAKEOVER-READY” message.

(238) When the TP **314b2** receives the takeover ready message, the TP **314b2** sends a Technical Failure notification to customers. At this stage, the TP **314b2** does not accept any new electronic messages including orders from anyone for a specific time, e.g., 10 seconds. The TP **314b2** then initiates a “Reset” process to reinstate current order books states based on the backup state, process cancel on disconnect and sends cancellations, and starts accepting new electronic messages with new orders for processing as the new primary instance. The “Reset” process enables recovery in case the connection with customers is lost. In this case, customers would miss some messages and their order books may not be synchronized. The “Reset” process includes cleanly resetting the order books and reinstating the order books based on the backup state. Therefore, any stale messages from the customers are cleaned. When the “Reset” process is successful, the TP **314b2** notifies the TR **314b1** via a channel reset message.

(239) When the TR **314b1** receives the channel reset message, the TR **314b1** triggers a cancel to disconnect TCP connections, e.g., TM connections, that were connected with the transaction processing server **306a** designated as the old primary instance. The TP **314b2** processes the disconnection, cancels the orders before accepting new orders, and sends cancellation notices to the customers.

(240) The TP **314b2** completes the “Reset” and starts accepting new orders for processing.

(241) At this stage, the transaction processing server **306b** designated as the new primary instance becomes the primary instance. Each of the other transaction processing servers designated as a backup instance continues to operate by receiving incoming requests from the new primary instance after a synchronization process with the new primary instance.

(242) FIG. **14** illustrates an example flowchart **1400** of the operation of the fault tolerance system **300** shown in FIGS. **3** and **5**. In one embodiment, the operation of the system **300** includes: providing a plurality of transaction processing servers **306** coupled with each other via an electronic communications network **220**, each of the plurality of transaction processing servers **306** comprising a processor (**1402**); implementing, by the processor of each of the plurality of transaction processing servers **306**, a plurality of components **314** coupled with each other, each of the plurality of components **314** characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof, wherein one of the plurality of transaction processing servers is designated as a primary instance **306a** and each of the other transaction processing servers is designated as a backup instance **306b**, wherein each of the plurality of transaction processing servers **306** is operative to receive an incoming request for a transaction, process the received request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and previously received request and generate a result indicative thereof, the transaction processing server designated as the primary instance **306a** at least commences processing of the incoming request prior to each of the transaction processing servers designated as the backup instance **306b**, wherein the one or more of the plurality of components **314a** of the transaction processing server designated as the primary instance **306a** are each configured to communicate an electronic result message indicative of a result of an operation thereof to the corresponding component of each of the other transaction processing servers designated as the backup instance **306b** (**1404**); determining, by a plurality of fault tolerance processors **316**, each of the plurality of fault tolerance processors coupled with one transaction processing server of the plurality of transaction processing servers, that a failure has occurred in the transaction processing server designated as the primary instance **306a** (**1406**); causing, by the plurality of fault tolerance processors **316**, one of the others of the plurality of transaction processing servers to be designated as a new primary instance (**1408**); synchronizing, by the transaction processing server which is designated as the new primary instance, the current state of each of the plurality of components thereof based on the received electronic messages indicative of a result of the operation the corresponding component of the primary instance (**1410**); and processing, by the transaction processing server, any incoming requests not processed by the primary instance before commencing processing of new incoming requests (**1412**).

(243) In one embodiment, upon receipt of the electronic result message indicative of the result of the operation of the one or more of the plurality of components of the primary instance, the corresponding component of each of the other transaction processing servers designated as the backup instance compares a result of the operation thereof with the result indicated by the received electronic result message to determine whether one of the plurality of components of the transaction processing server designated as the primary instance or the corresponding component of each of the other transaction processing servers designated as the backup instance is operating correctly.

(244) In one embodiment, the operation of the system **300** further includes transmitting, by only the transaction processing server designated as the primary instance **306a**, a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and previously received request to an external recipient.

(245) In one embodiment, the operation of the system **300** further includes receiving periodically, by each of the plurality of fault tolerant processors **316**, a communication from the others of the

plurality of fault tolerant processors, wherein a failure to receive the communication from one of the plurality of fault tolerant processors is deemed to indicate a failure of the transaction processing server coupled therewith.

(246) In one embodiment, each of the plurality of components **314** of each of the plurality of transaction processing servers are further coupled with the fault tolerance processor **316** coupled with that transaction processing server **306** such that the fault tolerance processor **316** is able to detect a failure of any of the plurality of components **314**.

(247) In one embodiment, the state of each of the plurality of components **314b** of each of the other transaction processing servers designated as the backup instance **306b** corresponds to the state of each of the plurality of components **314a** of the transaction processing server designated as the primary instance **306a** subsequent to the processing of an electronic transaction message preceding the electronic transaction message currently being processed by the transaction processing server designated as the primary instance **306a**.

(248) In one embodiment, the operation of the system **300** further includes establishing, by the plurality of fault tolerance processors **316**, upon the plurality of transaction processing servers **306** starting up, which of the plurality of transaction processing servers is designated as the primary instance **306a**.

(249) In one embodiment, each of the other transaction processing servers designated as the backup instance **306b** includes a queue configured to store previously received electronic messages from the transaction processing server designated as the primary instance **306a**. In one embodiment, the transaction processing server which is designated as the new primary instance is configured to process all of the electronic messages stored in the queue before becoming the new primary instance.

(250) In one embodiment, the electronic result message indicative of the result of the operation of the transaction processing server designated as the primary instance **306a** includes a primary result output. The electronic result message indicative of the result of the operation of each of the other transaction processing servers designated as the backup instance includes a backup result output.

(251) In one embodiment, at least a subset of the plurality of components **314** of the transaction processing server designated as the new primary instance are configured to complete a reconciliation of the backup result output with the primary result output.

(252) In one embodiment, at least a subset of the plurality of components **314** of the transaction processing server designated as the new primary instance are configured to identify a backup result output that does not have a respective primary instance result output and publish an unmatched result output via a shared memory architecture.

(253) In one embodiment, when the fault tolerance processor of the plurality of fault tolerance processors coupled with the primary instance **306a** determines that the failure has occurred in the primary instance, causing, by each of the plurality of fault tolerance processors, the primary instance **306a** to at least de-couple from the others of the plurality of transaction processing servers.

(254) In one embodiment, a first component of the plurality of components comprises a TR for receiving the electronic message via the electronic communications network, the electronic message including data indicative of an order to trade. A second component of the plurality of components comprises a TP for matching the received order to trade with one or more of previously received unmatched orders which are counter thereto. A third component of the plurality of components comprises an RG for generating and publishing a market data message.

(255) FIG. **15** depicts the state transitions of the FTs including startup, designation of primary and backup roles, failure detection, and primary take over.

(256) In one embodiment, an operation **1500** of the system **300** includes: initializing, by a FT **316a**; coupling, by the FT **316a**, with a candidate instance **306a** including a plurality components **314a**; waiting, by a FT **316a** for the plurality of components **314a** to join or register in the group; voting, by the FTs **316**, for a role including primary or backup to be assigned to the candidate instance

306a; designating the instance **306a** as a primary candidate **306a**; performing a health check of the primary candidate **306a**; initializing, by the primary candidate instance **306a**, a primary specific state; finalizing, by the FT **316a**, the designation of the instance **306** as a primary instance, waiting for the plurality of components **314a** to be ready, designating the primary candidate as Primary ready; and promoting the primary candidate to a primary instance **306a**.

(257) In one embodiment, the operation **1500** of the system **300** includes: initializing, by a FT **316b**; coupling, by the FT **316b**, with a candidate instance **306b** including a plurality components **314b**; waiting, by a FT **316b** for the plurality of components **314b** to join or register in the group; voting, by the FTs **316**, for a role including primary or backup to be assigned to the candidate instance **306b**; designating the instance **306b** as a backup candidate **306b**; performing a health check of the primary candidate **306b**; initializing, by the backup candidate instance **306b**, a backup specific state; finalizing the initializing of the backup candidate instance **306b** as a backup instance; waiting for the plurality of components **314b** to be ready; designating the backup candidate instance as Backup ready; and promoting the backup candidate instance to a backup instance **306b**.

(258) In one embodiment, the operation **1500** further includes, detecting, by any of the plurality of FTs, a failure in the primary instance **306a**; designating, by the FTs the backup instance **306b** as the new primary instance; processing, by the backup instance **306b**, all transaction messages to synchronize the state; initiating the takeover; taking over as a new primary instance; finalizing the takeover as a new primary instance; becoming a new primary instance; and processing new transaction messages.

CONCLUSION

(259) The illustrations of the embodiments described herein are intended to provide a general understanding of the structure of the various embodiments. The illustrations are not intended to serve as a complete description of all of the elements and features of apparatus and systems that utilize the structures or methods described herein. Many other embodiments may be apparent to those of skill in the art upon reviewing the disclosure. Other embodiments may be utilized and derived from the disclosure, such that structural and logical substitutions and changes may be made without departing from the scope of the disclosure. Additionally, the illustrations are merely representational and may not be drawn to scale. Certain proportions within the illustrations may be exaggerated, while other proportions may be minimized. Accordingly, the disclosure and the figures are to be regarded as illustrative rather than restrictive.

(260) While this specification contains many specifics, these should not be construed as limitations on the scope of the invention or of what may be claimed, but rather as descriptions of features specific to particular embodiments of the invention. Certain features that are described in this specification in the context of separate embodiments can also be implemented in combination in a single embodiment. Conversely, various features that are described in the context of a single embodiment can also be implemented in multiple embodiments separately or in any suitable sub-combination. Moreover, although features may be described as acting in certain combinations and even initially claimed as such, one or more features from a claimed combination can in some cases be excised from the combination, and the claimed combination may be directed to a sub-combination or variation of a sub-combination.

(261) Similarly, while operations are depicted in the drawings and described herein in a particular order, this should not be understood as requiring that such operations be performed in the particular order shown or in sequential order, or that all illustrated operations be performed, to achieve desirable results. In certain circumstances, multitasking and parallel processing may be advantageous. Moreover, the separation of various system components in the described embodiments should not be understood as requiring such separation in all embodiments, and it should be understood that the described program components and systems can generally be integrated together in a single software product or packaged into multiple software products.

(262) One or more embodiments of the disclosure may be referred to herein, individually and/or

collectively, by the term “invention” merely for convenience and without intending to voluntarily limit the scope of this application to any particular invention or inventive concept. Moreover, although specific embodiments have been illustrated and described herein, it should be appreciated that any subsequent arrangement designed to achieve the same or similar purpose may be substituted for the specific embodiments shown. This disclosure is intended to cover any and all subsequent adaptations or variations of various embodiments. Combinations of the above embodiments, and other embodiments not specifically described herein, will be apparent to those of skill in the art upon reviewing the description.

(263) The Abstract of the Disclosure is provided to comply with 37 C.F.R. § 1.72(b) and is submitted with the understanding that it will not be used to interpret or limit the scope or meaning of the claims. In addition, in the foregoing Detailed Description, various features may be grouped together or described in a single embodiment for the purpose of streamlining the disclosure. This disclosure is not to be interpreted as reflecting an intention that the claimed embodiments require more features than are expressly recited in each claim. Rather, as the following claims reflect, inventive subject matter may be directed to less than all of the features of any of the disclosed embodiments. Thus, the following claims are incorporated into the Detailed Description, with each claim standing on its own as defining separately claimed subject matter.

(264) It is therefore intended that the foregoing detailed description be regarded as illustrative rather than limiting, and that it be understood that it is the following claims, including all equivalents, that are intended to define the spirit and scope of this invention.

Claims

1. A system comprising: a plurality of transaction processing servers coupled with each other via an electronic communications network, each of the plurality of transaction processing servers comprising a processor which implements a plurality of components coupled with each other, each of the plurality of components characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof, wherein one of the plurality of transaction processing servers is designated as a primary instance and each other of the plurality of transaction processing servers is designated as a backup instance, wherein each of the plurality of transaction processing servers is operative to receive an incoming request for a transaction, process the received incoming request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and the previously received request and generate a result indicative thereof, the transaction processing server designated as the primary instance at least commencing processing of the incoming request prior to each of the other of the plurality of transaction processing servers designated as the backup instance, wherein one or more of the plurality of components of the transaction processing server designated as the primary instance are each configured to communicate the electronic result message indicative of the result of the operation thereof to a corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance; and a plurality of fault tolerance processors, each of the plurality of transaction processing servers coupled with one of the plurality of fault tolerance processors, each of the plurality of fault tolerance processors configured to determine that a failure has occurred in the transaction processing server designated as the primary instance and, subsequent thereto, cause one of the other of the plurality of transaction processing servers to be designated as a new primary instance, wherein the one of other of the plurality of transaction processing servers which is designated as the new primary instance is configured to synchronize the current state of each of the plurality of components thereof based on received electronic messages indicative of the result of the operation of the corresponding component of the primary

instance, and process any incoming requests not processed by the primary instance before commencing processing of new incoming requests.

2. The system of claim 1, wherein upon receipt of the electronic result message indicative of the result of the operation of the one or more of the plurality of components of the primary instance, the corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance compares a result of the operation thereof with the result indicated by the received electronic result message to determine whether one of the plurality of components of the one of the plurality of transaction processing servers designated as the primary instance or the corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance is operating correctly.

3. The system of claim 1, wherein only the one of the plurality of transaction processing servers designated as the primary instance is configured to transmit a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and the previously received request to a recipient external to the system.

4. The system of claim 1, wherein each of the plurality of fault tolerant processors periodically receives a communication from others of the plurality of fault tolerant processors, wherein a failure to receive the communication from one of the plurality of fault tolerant processors is deemed to indicate a failure of the transaction processing server coupled therewith.

5. The system of claim 1, wherein each of the plurality of components of each of the plurality of transaction processing servers are further coupled with the one of the plurality of fault tolerance processors coupled with that transaction processing server such that the one of the plurality of fault tolerance processors is able to detect a failure of any of the plurality of components.

6. The system of claim 1, wherein the current state of each of the plurality of components of each of the other of the plurality of transaction processing servers designated as the backup instance corresponds to a state of each of the plurality of components of the one of the plurality of transaction processing servers designated as the primary instance subsequent to the processing of an electronic transaction message preceding the electronic transaction message currently being processed by the one of the plurality of transaction processing servers designated as the primary instance.

7. The system of claim 1, wherein the plurality of fault tolerance processors, establish, upon the plurality of transaction processing servers starting up, which of the plurality of transaction processing servers is designated as the primary instance.

8. The system of claim 1, wherein each of the other of the plurality of transaction processing servers designated as the backup instance includes a queue configured to store previously received electronic messages from the one of the plurality of transaction processing servers designated as the primary instance, and wherein the one of the plurality of transaction processing servers which is designated as the new primary instance is configured to process all of the previously received electronic messages stored in the queue before becoming the new primary instance.

9. The system of claim 8, wherein the electronic message indicative of the result of the operation of the one of the plurality of transaction processing servers designated as the primary instance includes a primary result output, and wherein the electronic message indicative of the result of the operation of each of the other of the plurality of transaction processing servers designated as the backup instance includes a backup result output.

10. The system of claim 9, wherein at least a subset of the plurality of components of the one of other of the plurality of transaction processing servers designated as the new primary instance are configured to complete a reconciliation of the backup result output with the primary result output.

11. The system of claim 9, wherein at least a subset of the plurality of components of the one of other of the plurality of transaction processing servers designated as the new primary instance are configured to identify a backup result output that does not have a respective primary result output and publish an unmatched result output via a shared memory architecture.

12. The system of claim 1, wherein when the one of the plurality of fault tolerance processors coupled with the primary instance determines that the failure has occurred in the primary instance, each of the plurality of fault tolerance processors is configured to cause the primary instance to at least de-couple from the other of the plurality of transaction processing servers.

13. The system of claim 1, wherein a first component of the plurality of components comprises a transaction receiver processor for receiving the electronic message via the electronic communications network, the electronic message including data indicative of an order to trade, wherein a second component of the plurality of components comprises a transaction processor for matching the received order to trade with one or more of previously received unmatched orders which are counter thereto, and wherein a third component of the plurality of components comprises a result generator processor for generating and publishing a market data message.

14. A computer implemented method comprising: providing a plurality of transaction processing servers coupled with each other via an electronic communications network each of the plurality of transaction processing servers comprising a processor; implementing, by the processor of each of the plurality of transaction processing servers, a plurality of components coupled with each other, each of the plurality of components characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof, wherein one of the plurality of transaction processing servers is designated as a primary instance and each other of the plurality of transaction processing servers is designated as a backup instance, wherein each of the plurality of transaction processing servers is operative to receive an incoming request for a transaction, process the received incoming request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and the previously received request and generate a result indicative thereof, the transaction processing server designated as the primary instance at least commencing processing of the incoming request prior to each of the other of the plurality of transaction processing servers designated as the backup instance, wherein one or more of the plurality of components of the transaction processing server designated as the primary instance are each configured to communicate the electronic result message indicative of the result of the operation thereof to a corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance; determining, by a plurality of fault tolerance processors, each of the plurality of fault tolerance processors coupled with one transaction processing server of the plurality of transaction processing servers, that a failure has occurred in the transaction processing server designated as the primary instance; causing, by the plurality of fault tolerance processors, one of the other of the plurality of transaction processing servers to be designated as a new primary instance; synchronizing, by the transaction processing server which is designated as the new primary instance, the current state of each of the plurality of components thereof based on received electronic messages indicative of the result of the operation of the corresponding component of the primary instance; and processing, by the transaction processing server, any incoming requests not processed by the primary instance before commencing processing of new incoming requests.

15. The method of claim 14, wherein upon receipt of the electronic result message indicative of the result of the operation of the one or more of the plurality of components of the primary instance, the corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance compares a result of the operation thereof with the result indicated by the received electronic result message to determine whether one of the plurality of components of the one of the plurality of transaction processing servers designated as the primary instance or the corresponding component of the other of the plurality of transaction processing servers designated as the backup instance is operating correctly.

16. The method of claim 14, further comprising: transmitting, by only the one of the plurality of

transaction processing servers designated as the primary instance, a data message indicative of the generated result of the attempt to satisfy one or both of the incoming request and the previously received request to an external recipient.

17. The method of claim 14, further comprising: receiving periodically, by each of the plurality of fault tolerant processors, a communication from others of the plurality of fault tolerant processors, wherein a failure to receive the communication from one of the plurality of fault tolerant processors is deemed to indicate a failure of the transaction processing server coupled therewith.

18. The method of claim 14, wherein each of the plurality of components of each of the plurality of transaction processing servers are further coupled with the one of the plurality of fault tolerance processors coupled with that transaction processing server such that the one of the plurality of fault tolerance processors is able to detect a failure of any of the plurality of components.

19. The method of claim 14, wherein the current state of each of the plurality of components of each of the other of the plurality of transaction processing servers designated as the backup instance corresponds to a state of each of the plurality of components of the one of the plurality of transaction processing servers designated as the primary instance subsequent to the processing of an electronic transaction message preceding the electronic transaction message currently being processed by the one of the plurality of transaction processing servers designated as the primary instance.

20. The method of claim 14, further comprising: establishing, by the plurality of fault tolerance processors, upon the plurality of transaction processing servers starting up, which of the plurality of transaction processing servers is designated as the primary instance.

21. The method of claim 14, wherein each of the other of the plurality of transaction processing servers designated as the backup instance includes a queue configured to store previously received electronic messages from the one of the plurality of transaction processing servers designated as the primary instance, and wherein the one of the other of the plurality of transaction processing servers which is designated as the new primary instance is configured to process all of the previously received electronic messages stored in the queue before becoming the new primary instance.

22. The method of claim 21, wherein the electronic message indicative of the result of the operation of the one of the plurality of transaction processing servers designated as the primary instance includes a primary result output, and wherein the electronic message indicative of the result of the operation of each of the other of the plurality of transaction processing servers designated as the backup instance includes a backup result output.

23. The method of claim 22, wherein at least a subset of the plurality of components of the one of other of the plurality of transaction processing servers designated as the new primary instance are configured to complete a reconciliation of the backup result output with the primary result output.

24. The method of claim 21, wherein at least a subset of the plurality of components of the one of the other of the plurality of transaction processing servers designated as the new primary instance are configured to identify a backup result output that does not have a respective primary result output and publish an unmatched result output via a shared memory architecture.

25. The method of claim 14, further comprising: when the one of the plurality of fault tolerance processors coupled with the primary instance determines that the failure has occurred in the primary instance, causing, by each of the plurality of fault tolerance processors, the primary instance to at least de-couple from the other of the plurality of transaction processing servers.

26. The method of claim 14, wherein a first component of the plurality of components comprises a transaction receiver processor for receiving the electronic message via the electronic communications network, the electronic message including data indicative of an order to trade, wherein a second component of the plurality of components comprises a transaction processor for matching the received order to trade with one or more of previously received unmatched orders which are counter thereto, and wherein a third component of the plurality of components comprises a result generator processor for generating and publishing a market data message.

27. A system comprising: means for providing a plurality of transaction processing servers coupled with each other via an electronic communications network; means for implementing a plurality of components of each of the plurality of transaction processing servers, coupled with each other, each of the plurality of components characterized by a current state and configured to perform one or more of receiving an electronic message, processing the received electronic message and generating an electronic result message indicative of a result of an operation which may alter the current state thereof, wherein one of the plurality of transaction processing servers is designated as a primary instance and each other of the plurality of transaction processing servers is designated as a backup instance, wherein each of the plurality of transaction processing servers is operative to receive an incoming request for a transaction, process the received incoming request to identify a previously received request for a transaction counter thereto in an attempt to satisfy one or both of the incoming request and the previously received request and generate a result indicative thereof, the transaction processing server designated as the primary instance at least commencing processing of the incoming request prior to each of the other of the plurality of transaction processing servers designated as the backup instance, wherein one or more of the plurality of components of the transaction processing server designated as the primary instance are each configured to communicate the electronic result message indicative of the result of the operation thereof to a corresponding component of each of the other of the plurality of transaction processing servers designated as the backup instance; means for determining that a failure has occurred in the transaction processing server designated as the primary instance; means for causing one of the other of the plurality of transaction processing servers to be designated as a new primary instance; means for synchronizing the current state of each of the plurality of components thereof based on received electronic messages indicative of the result of the operation of the corresponding component of the primary instance; and means for processing any incoming requests not processed by the primary instance before commencing processing of new incoming requests.
