

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417168
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Galvin; Brian

System and method for edge based multi-modal homomorphic compression

Abstract

A system and method for compressing and restoring multi-modal data utilizing a variational autoencoder to enable homomorphic compression techniques is disclosed. Multi-modal input data, comprising at least two different data types, is compressed into a unified latent space using an encoder network of a multi-modal variational autoencoder. Homomorphic operations are performed on compressed data in the latent space. The latent space compressed data is decompressed using a decoder network of the multi-modal variational autoencoder. The system utilizes modality-specific layers and cross-modal attention mechanisms to effectively process diverse data types. The homomorphic operations enable performing computations while the data is in a compressed form, preserving results of those operations in the decompressed output. This approach allows for efficient storage, transmission, and analysis of multi-modal data while maintaining privacy and data integrity across different modalities.

Inventors:	Galvin; Brian (Silverdale, WA)
Applicant:	AtomBeam Technologies Inc. (Moraga, CA)
Family ID:	1000008820416
Assignee:	ATOMBEAM TECHNOLOGIES INC. (Moraga, CA)
Appl. No.:	19/048882
Filed:	February 08, 2025

Prior Publication Data

Document Identifier	Publication Date
US 20250190339 A1	Jun. 12, 2025

Related U.S. Application Data

continuation-in-part parent-doc US 18890774 20240920 US 12225105 child-doc US 19048882
continuation-in-part parent-doc US 18657683 20240507 US 12159216 20241203 child-doc US 18890774
continuation-in-part parent-doc US 18648340 20240427 US 12166507 20241210 child-doc US 18657683
continuation-in-part parent-doc US 18427716 20240130 US 12093972 20240917 child-doc US 18648340
continuation-in-part parent-doc US 18410980 20240111 US 12068761 20240820 child-doc US 18427716
continuation-in-part parent-doc US 18537728 20231212 US 12058333 20240806 child-doc US 18410980

Publication Classification

Int. Cl.: H03M7/00 (20060101); G06F12/02 (20060101); H04L9/00 (20220101)

U.S. Cl.:

CPC G06F12/0223 (20130101); H04L9/008 (20130101);

Field of Classification Search

CPC: H03M (7/3062); H04N (19/80)

USPC: 341/50; 341/51

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
7629922	12/2008	Winstead et al.	N/A	N/A
7876257	12/2010	Vetro et al.	N/A	N/A
10956787	12/2020	Rothberg et al.	N/A	N/A
11237549	12/2021	Kane	N/A	G05B 19/418
11656353	12/2022	Li et al.	N/A	N/A
11915690	12/2023	Chang et al.	N/A	N/A
2020/0272605	12/2019	More	N/A	N/A
2020/0404340	12/2019	Xu et al.	N/A	N/A
2022/0246158	12/2021	Nam et al.	N/A	N/A

FOREIGN PATENT DOCUMENTS

Patent No.	Application Date	Country	CPC
118677928	12/2023	CN	G08B 21/18

OTHER PUBLICATIONS

Cai, Yuanhao et al., “Mask-guided Spectral-wise Transformer for Efficient Hyperspectral Image Reconstruction”, Conference on Computer Vision and Pattern Recognition, pp. 17502-17511, 2022.
cited by applicant

Cai, Yuanhao et al., “MST++: Multi-stage Spectral-wise Transformer for Efficient Spectral

Reconstruction”, Conference on Computer Vision and Pattern Recognition, 2022. cited by applicant

He, Kaiming et al., “Deep Residual Learning for Image Recognition”, Conference on Computer Vision and Pattern Recognition, pp. 770-778, 2016. cited by applicant

Primary Examiner: Jeanglaude; Jean B

Attorney, Agent or Firm: Galvin Patent Law LLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) Priority is claimed in the application data sheet to the following patents or patent applications, each of which is expressly incorporated herein by reference in its entirety: Ser. No. 18/890,774 Ser. No. 18/657,683 Ser. No. 18/648,340 Ser. No. 18/427,716 Ser. No. 18/410,980 Ser. No. 18/537,728

BACKGROUND OF THE INVENTION

Field of the Art

(2) The present invention relates to the field of deep learning and data compression. More specifically, the invention pertains to systems and methods that perform homomorphic compression and decompression.

Discussion of the State of the Art

(3) In recent years, deep learning approaches have shown promising results in multi-modal data compression and restoration. Autoencoders, a type of neural network architecture, have emerged as a powerful tool for learning compact representations of diverse data types. Traditional autoencoders consist of an encoder network that maps input data to a lower-dimensional latent space and a decoder network that reconstructs the original data from the latent representation. However, these have been primarily focused on single modality data.

(4) Multi-modal autoencoders have been proposed to handle diverse data types simultaneously, such as text, images, audio, and sensor data. These architectures extend the traditional autoencoder concept to learn joint representations of multiple data modalities. Multi-layer multi-modal autoencoders, also known as stacked multi-modal autoencoders or deep multi-modal autoencoders, have been developed to learn hierarchical representations of multi-modal data. These architectures stack multiple layers of encoders and decoders, often with modality-specific and shared layers, allowing for more complex and abstract feature learning across different data types. Multi-layer multi-modal autoencoders have been successfully applied in various domains, such as cross-modal retrieval, multi-modal fusion for sentiment analysis, and multi-sensor data compression in IoT applications.

(5) Multi-modal autoencoders are becoming an essential component of many advanced machine learning systems, particularly those dealing with diverse data sources. Machine learning is becoming more prevalent for a variety of important tasks that often involve multi-modal data, such as identifying trends and patterns in large volumes of heterogeneous data, performing automated tasks and processes that require integration of different data types, and enabling improved customer experiences and personalization based on multi-faceted user data. The ability to efficiently compress, process, and reconstruct multi-modal data is important in fields such as healthcare (combining medical images, patient records, and sensor data), autonomous vehicles (integrating visual, lidar, and GPS data), and smart cities (processing diverse IoT sensor data). As technology continues to improve, machine learning and artificial intelligence (AI) systems capable of handling multi-modal data are likely to find new uses and applications, particularly in scenarios where the

integration of diverse data types can provide more comprehensive insights and capabilities.

SUMMARY OF THE INVENTION

(6) Disclosed embodiments provide a system and method for compressing and restoring multi-modal data using variational autoencoders on an edge device. The general architecture of an autoencoder includes an encoder, decoder, and bottleneck layer. The encoder section includes an input layer that accepts input data. Hidden layers progressively reduce the dimensionality of the input, capturing important features and patterns. The bottleneck layer (latent space) is the final hidden layer, where the dimensionality is significantly reduced. This layer represents the compressed encoding of the input data. In the decoder section, the encoded representation is received from the bottleneck layer and expanded back to the dimensionality of the original input. One or more hidden layers progressively increase the dimensionality and aim to reconstruct the original input. The output layer produces the reconstructed output, may resemble the input data. The loss function used during training is typically a reconstruction loss, measuring the difference between the input and the reconstructed output. In embodiments, the loss function can include mean squared error (MSE) for continuous data or binary cross-entropy for binary data. During training, the autoencoder learns to minimize the reconstruction loss, forcing the network to capture the most important features of the input data in the bottleneck layer.

(7) The architecture of the multi-modal variational autoencoder can include modality-specific layers for processing different data types, shared layers for integrating information across modalities, and activation functions. The multi-modal variational autoencoder can provide a unified latent space that is continuously differentiable. The continuously differentiable unified latent space can enable homomorphic operations, thereby providing homomorphic compression capabilities. Homomorphic operations allow computations to be performed on encrypted data without decrypting it first. Thus, these operations maintain the mathematical structure of the data even after transformation. The homomorphic properties enable important features of computer systems in areas such as security and privacy. For example, operations can be performed on sensitive information in its encrypted form, thereby enabling multiple parties to operate on the data without having the unencrypted contents revealed. This feature has applications for the handling of medical records, financial records, and/or any other records that contain sensitive information.

(8) According to a preferred embodiment, a computer system comprising a hardware memory, wherein the computer system is configured to execute software instructions stored on nontransitory machine-readable storage media that: preprocess raw multi-modal data at the edge computing device to generate a plurality of input data sets; compress the plurality of input data sets into a plurality of compressed data sets using an encoder within a multi-modal variational autoencoder at the edge computing device; transmit the compressed data sets from the edge computing device to the cloud computing device; and decompress the plurality of compressed data sets using a decoder located within the cloud computing device to obtain a plurality of reconstructed data sets, is disclosed.

(9) According to another preferred embodiment, A method for compressing and restoring multi-modal data, comprising steps of: preprocessing raw multi-modal data at the edge computing device to generate a plurality of input data sets; compressing the plurality of input data sets into a plurality of compressed data sets using an encoder within a multi-modal variational autoencoder at the edge computing device; transmitting the compressed data sets from the edge computing device to the cloud computing device; and decompressing the plurality of compressed data sets using a decoder located within the cloud computing device to obtain a plurality of reconstructed data sets, is disclosed.

(10) According to an aspect of an embodiment, the encoder and decoder of the multi-modal variational autoencoder comprise modality-specific layers for processing different data types, shared layers for integrating information across modalities, and activation functions.

(11) According to an aspect of an embodiment, the modality-specific layers comprise at least one of

convolutional layers for image data, recurrent layers for sequential data, and fully connected layers for numerical data.

(12) According to an aspect of an embodiment, the multi-modal data comprises at least two of: image data, audio data, text data, time-series data, and sensor data.

(13) According to an aspect of an embodiment, the multi-modal variational autoencoder is configured to perform partially homomorphic compression in a unified latent space representing all data modalities.

(14) According to an aspect of an embodiment, the multi-modal variational autoencoder comprises a unified latent space, and wherein the plurality of programming instructions further includes instructions that, when executing on the processor, cause the multi-modal variational autoencoder to perform one or more linear operations in the unified latent space.

(15) According to an aspect of an embodiment, the multi-modal variational autoencoder utilizes cross-modal attention mechanisms to capture relationships between different data types.

(16) According to an aspect of an embodiment, the method further comprises aligning and synchronizing the different data types during preprocessing and reconstruction.

(17) According to an aspect of an embodiment, the multi-modal variational autoencoder is trained using a loss function that accounts for reconstruction quality across all data modalities.

(18) According to an aspect of an embodiment, the system further comprises a multi-modal correlation network, and wherein the output of the decompressed output organizer is input to the multi-modal correlation network, and wherein the output of the multi-modal correlation network provides the reconstructed multi-modal output.

Description

BRIEF DESCRIPTION OF THE DRAWING FIGURES

(1) FIG. 1 is a block diagram illustrating an exemplary system architecture for compressing and restoring data using multi-level autoencoders and correlation networks.

(2) FIG. 2 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder network.

(3) FIG. 3 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network.

(4) FIG. 4 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder training system.

(5) FIG. 5 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, correlation network training system.

(6) FIG. 6 is a flow diagram illustrating an exemplary method for compressing a data input using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

(7) FIG. 7 is a flow diagram illustrating an exemplary method for decompressing a compressed data input using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

(8) FIG. 8 is a block diagram illustrating an exemplary system architecture for compressing and restoring IoT sensor data using a system for compressing and restoring data using multi-level autoencoders and correlation networks.

(9) FIG. 9 is a flow diagram illustrating an exemplary method for compressing and

(10) decompressing IoT sensor data using a system for compressing and restoring data using multi-

level autoencoders and correlation networks.

(11) FIG. **10** is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, the decompressed output organizer.

(12) FIG. **11** is a flow diagram illustrating an exemplary method for organizing restored, decompressed data sets after correlation network processing.

(13) FIG. **12** illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part.

(14) FIG. **13** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using variational autoencoders, according to an embodiment.

(15) FIG. **14** is a block diagram of a variational autoencoder used in an embodiment.

(16) FIG. **15** is a flow diagram illustrating an exemplary method for compressing a data input using a system for compressing and restoring data using variational autoencoders, according to an embodiment.

(17) FIG. **16** is a flow diagram illustrating an exemplary method for decompressing a compressed data input using a system for compressing and restoring data using variational autoencoders, according to an embodiment.

(18) FIG. **17** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using multi-level autoencoders and correlation networks on a plurality of edge devices.

(19) FIG. **18** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using variational autoencoders, according to an embodiment.

(20) FIG. **19** is a flow diagram illustrating an exemplary method for compressing a multi-modal data input on an edge device using a system for compressing and restoring data using variational autoencoders, according to an embodiment.

(21) FIG. **20** is a flow diagram illustrating an exemplary method for decompressing compressed multi-modal data input on using a system for compressing and restoring data using variational autoencoders stores in a cloud environment, according to an embodiment.

DETAILED DESCRIPTION OF THE INVENTION

(22) The inventor has conceived and reduced to practice a system and method for compressing and restoring multi-modal data on an edge device. The system comprises a computing device with at least a memory and a processor, and a plurality of programming instructions stored in the memory and operable on the processor. When executed, these instructions cause the computing device to perform several key operations.

(23) First, the system preprocesses raw multi-modal data to generate a plurality of input data sets. Importantly, this multi-modal data comprises at least two different data types, allowing the system to handle diverse forms of information simultaneously. The system then compresses these input data sets into a plurality of compressed data sets using an encoder within a multi-modal variational autoencoder. This compression step is helpful for reducing data size while preserving essential information across all modalities.

(24) Following compression, the system decompresses the compressed data sets using a decoder located within the same multi-modal variational autoencoder, obtaining a plurality of reduced output data sets. Finally, these reduced output data sets are processed through a decompressed output organizer, generating a reconstructed multi-modal output that closely resembles the original input data.

(25) The multi-modal variational autoencoder at the heart of this system incorporates several sophisticated features. Its encoder and decoder comprise modality-specific layers for processing different data types, shared layers for integrating information across modalities, and activation functions. The modality-specific layers include convolutional layers for image data, recurrent layers for sequential data, and fully connected layers for numerical data, allowing the system to

handle a wide variety of data types effectively.

(26) One of the key innovations of this system is its ability to perform partially homomorphic compression in a unified latent space representing all data modalities. This allows for certain operations to be performed on the compressed data without fully decompressing it, offering significant advantages in terms of data privacy and computational efficiency. The system performs one or more linear operations in this unified latent space, further extending its capabilities.

(27) The multi-modal variational autoencoder also utilizes cross-modal attention mechanisms to capture relationships between different data types, ensuring that important inter-modal correlations are preserved during the compression and decompression process.

(28) In some implementations, the system includes a multi-modal correlation network. In this case, the output of the decompressed output organizer serves as input to this network, which then provides the final reconstructed multi-modal output. This additional step further enhances the quality of the reconstructed data by leveraging correlations across different modalities.

(29) The inventor envisions numerous practical applications for this multi-modal homomorphic compression system. In a non-limiting use case example, the system is utilized in an industry where data privacy and security are paramount, such as healthcare or finance. For instance, a business with secure data needs, such as a hospital or a bank, uses the system to compress and analyze diverse customer or patient data. This data includes various types such as images, textual records, and time-series data from monitoring devices or transactions. The homomorphic properties of the compression allow for analysis of this data while maintaining privacy, facilitating large-scale studies or analyses while ensuring compliance with relevant data protection regulations.

(30) In another non-limiting use case example, the system is employed in a field requiring real-time processing of multi-modal sensor data, such as autonomous vehicles or advanced robotics. The system efficiently compresses and processes multi-modal sensor data in real-time, including visual data, distance measurement data, and location information. The compressed format allows for efficient data transmission to remote processing centers for further analysis and machine learning model updates, helpful for the continuous improvement of the system's performance and decision-making capabilities.

(31) In a further non-limiting use case example, the system is applied in a context involving multiple data sources and privacy concerns, such as smart city management or industrial IoT. The system compresses and analyzes data from various sensors measuring different environmental or operational factors, alongside video feeds from monitoring cameras. The homomorphic properties allow for data analysis and decision-making without fully decrypting sensitive information, thereby enhancing privacy for individuals or protecting proprietary information while still enabling effective management and optimization of operations.

(32) These use cases illustrate the wide-ranging potential of the invented system, showcasing its ability to handle diverse data types, preserve privacy, and enable efficient analysis across various industries and applications. The system's versatility makes it applicable in numerous scenarios where multi-modal data compression, privacy preservation, and efficient analysis are required.

(33) One or more different aspects may be described in the present application. Further, for one or more of the aspects described herein, numerous alternative arrangements may be described; it should be appreciated that these are presented for illustrative purposes only and are not limiting of the aspects contained herein or the claims presented herein in any way. One or more of the arrangements may be widely applicable to numerous aspects, as may be readily apparent from the disclosure. In general, arrangements are described in sufficient detail to enable those skilled in the art to practice one or more of the aspects, and it should be appreciated that other arrangements may be utilized and that structural, logical, software, electrical and other changes may be made without departing from the scope of the particular aspects. Particular features of one or more of the aspects described herein may be described with reference to one or more particular aspects or figures that form a part of the present disclosure, and in which are shown, by way of illustration, specific

arrangements of one or more of the aspects. It should be appreciated, however, that such features are not limited to usage in the one or more particular aspects or figures with reference to which they are described. The present disclosure is neither a literal description of all arrangements of one or more of the aspects nor a listing of features of one or more of the aspects that must be present in all arrangements.

(34) Headings of sections provided in this patent application and the title of this patent application are for convenience only, and are not to be taken as limiting the disclosure in any way.

(35) Devices that are in communication with each other need not be in continuous communication with each other, unless expressly specified otherwise. In addition, devices that are in communication with each other may communicate directly or indirectly through one or more communication means or intermediaries, logical or physical.

(36) A description of an aspect with several components in communication with each other does not imply that all such components are required. To the contrary, a variety of optional components may be described to illustrate a wide variety of possible aspects and in order to more fully illustrate one or more aspects. Similarly, although process steps, method steps, algorithms or the like may be described in a sequential order, such processes, methods and algorithms may generally be configured to work in alternate orders, unless specifically stated to the contrary. In other words, any sequence or order of steps that may be described in this patent application does not, in and of itself, indicate a requirement that the steps be performed in that order. The steps of described processes may be performed in any order practical. Further, some steps may be performed simultaneously despite being described or implied as occurring non-simultaneously (e.g., because one step is described after the other step). Moreover, the illustration of a process by its depiction in a drawing does not imply that the illustrated process is exclusive of other variations and modifications thereto, does not imply that the illustrated process or any of its steps are necessary to one or more of the aspects, and does not imply that the illustrated process is preferred. Also, steps are generally described once per aspect, but this does not mean they must occur once, or that they may only occur once each time a process, method, or algorithm is carried out or executed. Some steps may be omitted in some aspects or some occurrences, or some steps may be executed more than once in a given aspect or occurrence.

(37) When a single device or article is described herein, it will be readily apparent that more than one device or article may be used in place of a single device or article. Similarly, where more than one device or article is described herein, it will be readily apparent that a single device or article may be used in place of the more than one device or article. The functionality or the features of a device may be alternatively embodied by one or more other devices that are not explicitly described as having such functionality or features. Thus, other aspects need not include the device itself.

(38) Techniques and mechanisms described or referenced herein will sometimes be described in singular form for clarity. However, it should be appreciated that particular aspects may include multiple iterations of a technique or multiple instantiations of a mechanism unless noted otherwise. Process descriptions or blocks in figures should be understood as representing modules, segments, or portions of code which include one or more executable instructions for implementing specific logical functions or steps in the process. Alternate implementations are included within the scope of various aspects in which, for example, functions may be executed out of order from that shown or discussed, including substantially concurrently or in reverse order, depending on the functionality involved, as would be understood by those having ordinary skill in the art.

(39) Conceptual Architecture

(40) FIG. 1 is a block diagram illustrating an exemplary system architecture for compressing and restoring multi-modal data using multi-level autoencoders and correlation networks. In one embodiment, a system for compressing and restoring multi-modal data using multi-level autoencoders and correlation networks comprises a plurality of multi-modal data inputs **100**, a data

preprocessor **110**, a data normalizer **120**, a multi-layer multi-modal variational autoencoder network **130** which further comprises an encoder network **131** and a decoder network **132**, a plurality of compressed outputs **140**, plurality of decompressed multi-modal outputs **170**, a decompressed output organizer **190**, a multi-modal correlation network **160**, and a reconstructed multi-modal output **180**. The plurality of multi-modal data inputs **100** are representations of raw data from various sources, such as sensors, cameras, or databases, comprising at least two different data types. These data types may include, but are not limited to, image data, audio data, text data, time-series data, and sensor data. The raw multi-modal data can be in different formats, including but not limited to images, videos, audio, or structured data. The plurality of multi-modal data inputs **100** may be transferred to the data preprocessor **110** for further processing.

(41) The data preprocessor **110** applies various preprocessing techniques to the raw multi-modal data received from the data input **100**. These techniques may include data cleaning, noise reduction, artifact removal, format conversion, and importantly, aligning and synchronizing the different data types. The preprocessor **110** ensures that the multi-modal data is in a suitable format and quality for subsequent stages of the system, with different modalities properly aligned for joint processing.

(42) The compressed latent representation obtained from the encoder network **131** is the compressed output **140**. The compressed output **140** has a significantly reduced size compared to the original input data, enabling efficient storage and transmission. The compressed output **140** may be stored in a storage system. A storage system may include any suitable storage medium, such as a database, file system, or cloud storage. Storage systems allow for the efficient management and retrieval of the compressed data as needed. When the compressed data needs to be restored or reconstructed, it may be retrieved from the storage system and passed to the decoder network **132**. Additionally, the compressed data may be directly passed to either the decompression network **132**. The decoder network **132** is responsible for decoding the compressed latent representation back into the original data space by outputting a decompressed multi-modal output **170**. It consists of multiple layers of decoders that progressively increase the dimensionality of the data, reconstructing the original input.

(43) The decompressed multi-modal output **170** from the decoder network **132** may have some loss of information compared to the original input data due to the compression process. To further enhance the quality of the decompressed output, the system may incorporate a multi-modal correlation network **160**. The correlation network **160** leverages the correlations and patterns between different compressed inputs to restore the decompressed output more accurately. It learns to capture the relationships and dependencies within the data, allowing for better reconstruction and restoration of the original information. The multi-modal correlation network **160** takes the decompressed multi-modal outputs **170** as inputs. It analyzes the correlations and similarities between the data samples and uses this information to refine and enhance the decompressed output. The refined decompressed output from the multi-modal correlation network **160** is a reconstructed multi-modal output **180** of the system. The reconstructed multi-modal output **180** closely resembles the original input data, with minimal loss of information and improved quality compared to the output from the decoder network **132** alone.

(44) In one embodiment, the multi-modal correlation network **160** may receive inputs from a decompressed output organizer **190** which operates on the decompressed multi-modal outputs **170** obtained from the decoder network **132**. The decompressed output organizer **190** may organize the decompressed multi-modal outputs **170** into groups based on their cross-modal correlations and similarities.

(45) By grouping decompressed multi-modal outputs **170** based on similarities across different data types, the multi-modal correlation network **160** will more easily be able to identify correlations between decompressed outputs **170** both within and across modalities. The multi-modal correlation network **160** finds patterns and similarities between decompressed outputs **170** to develop a more

holistic reconstructed original multi-modal input. By priming the multi-modal correlation network **160** with already grouped, similar compressed multi-modal outputs **170**, the network will be able to generate even more reliable reconstructions. The multi-layer multi-modal variational autoencoder network **130** and the multi-modal correlation network **160** are trained using a large dataset of diverse multi-modal samples. The training process involves minimizing the reconstruction loss between the original multi-modal input data and the decompressed output **170**, accounting for reconstruction quality across all data modalities. The system learns to compress the multi-modal data efficiently while preserving the essential features and patterns across different data types.

(46) FIG. 2 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring multi-modal data using multi-level autoencoders and correlation networks, a multi-layer multi-modal variational autoencoder network. The multi-layer multi-modal variational autoencoder network comprises a multi-modal encoder network **131** and a multi-modal decoder network **132** that work together to encode and decode multi-modal data effectively. The multi-modal encoder network **131** and multi-modal decoder network **132** within the multi-layer multi-modal variational autoencoder network are comprised of a plurality of layers that contribute to the encoding and decoding process. These layers include, but are not limited to, modality-specific layers for processing different data types, shared layers for integrating information across modalities, and a unified bottleneck layer. Some embodiments also include functions that operate on information including but not limited to rectified linear unit functions, sigmoid functions, and skip connections.

(47) The modality-specific layers are responsible for extracting meaningful features from the different types of input data. For image data, convolutional layers apply convolutional operations using learnable filters to capture spatial patterns and hierarchical representations of the data. For sequential data such as text or time-series, recurrent layers may be used. For numerical data, fully connected layers can be employed. These layers can have different configurations to capture features at various scales and resolutions appropriate to each data type. Skip connections are employed to facilitate the flow of information across different layers of the autoencoder. Skip connections allow the output of a layer to be directly added to the output of a subsequent layer, enabling the network to learn residual mappings and mitigate the vanishing gradient problem. Skip connections help in preserving fine-grained details and improving the training stability of the multi-modal autoencoder.

(48) Pooling layers are used to downsample the feature maps generated by the modality-specific layers. They reduce the dimensions of the feature maps while retaining the most salient information. Common pooling operations include but are not limited to max pooling and average pooling. Pooling layers help in achieving translation invariance, reducing computational complexity, and controlling the receptive field of the autoencoder. Rectified Linear Unit (ReLU) functions introduce non-linearity into the autoencoder by applying a ReLU activation function element-wise to the output of the previous layer. ReLU functions help in capturing complex patterns and relationships in the multi-modal data by allowing the network to learn non-linear transformations. They also promote sparsity and alleviate the vanishing gradient problem. The unified bottleneck layer represents the most compressed representation of the multi-modal input data in a shared latent space. This layer has a significantly reduced dimensionality compared to the input and output layers of the autoencoder. It forces the network to learn a compact and meaningful encoding of the data, capturing the essential features across all modalities and discarding redundant information. In one embodiment, the multi-layer multi-modal variational autoencoder network is comprised of a plurality of the previously mentioned layers where the sequence and composition of the layers may vary depending on a user's preferences and goals. The unified bottleneck layer is where the compressed multi-modal output **140** is created in a shared latent space. Each layer previous to the bottleneck layer creates a more and more compressed version of the original multi-modal input. The layers after the bottleneck layer represent the decoder network **132** where a

plurality of layers operate on a compressed input to decompress a multi-modal data set. Decompression results in a version of the original multi-modal input which is largely similar but has some lost data from the transformations.

(49) FIG. 3 is a block diagram illustrating an exemplary architecture for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network. The correlation network **160** is designed to enhance the reconstruction of decompressed data by leveraging correlations and patterns within the data. The correlation network **160** may also be referred to as a neural upsampler. The correlation network **160** comprises a plurality of correlation network elements that work together to capture and utilize the correlations for improved data reconstruction. Each correlation network element within the correlation network **160** contributes to the correlation learning and data reconstruction process. These elements include, but are not limited to, convolutional layers, skip connections, pooling layers and activation functions such as but not limited to, rectified linear unit functions or sigmoid functions.

(50) The convolutional layers are responsible for extracting meaningful features from the input data. They apply convolutional operations using learnable filters to capture spatial patterns and hierarchical representations of the data. The convolutional layers can have different numbers of filters, kernel sizes, and strides to capture features at various scales and resolutions. Skip connections are employed to facilitate the flow of information across different layers of the autoencoder. Skip connections allow the output of a layer to be directly added to the output of a subsequent layer, enabling the network to learn residual mappings and mitigate the vanishing gradient problem. Skip connections help in preserving fine-grained details and improving the training stability of the autoencoder.

(51) Pooling layers are used to downsample the feature maps generated by the convolutional layers. They reduce the spatial dimensions of the feature maps while retaining the most salient information. Common pooling operations include but are not limited to max pooling and average pooling. Pooling layers help in achieving translation invariance, reducing computational complexity, and controlling the receptive field of the autoencoder. Rectified Linear Unit (ReLU) functions introduce non-linearity into the autoencoder by applying a ReLU activation function element-wise to the output of the previous layer. ReLU functions help in capturing complex patterns and relationships in the data by allowing the network to learn non-linear transformations. They also promote sparsity and alleviate the vanishing gradient problem.

(52) In one embodiment, the correlation network **160** may comprise an encoder **310**, a decoder **320**, an N number of correlated data sets **300**, an N number-channel wise transformer **330**, and an N number of restored data sets **340**. Additionally, the correlation network **160** may be comprised of a plurality of convolutional layers, pooling layers, and activation functions. In one embodiment, the correlation network **160** may be configured to receive N correlated data sets **300** where each correlated data set includes a plurality of decompressed data points. In one embodiment, the correlation network **160** may be configured to receive four correlated data sets as an input. The correlated data sets may have been organized by a decompressed output organizer **190** to maximize the similarities between the data points in each set. One data set, **300**, may include data points **300a**, **300b**, **300c**, through **300n**, where the decompressed output organizer **190** has determined the N number of data points are similar enough to be grouped together. The correlation network **160** may then receive and process full data sets at a time. In FIG. 3, the data is processed through an encoder **310** by passing through a convolutional layer, a pooling layer, and an activation function.

(53) Activation functions introduce non-linearity into the network, enabling it to learn and represent complex patterns and relationships in the data. Common activation functions include but are not limited to sigmoid, tan h, ReLU (Rectified Linear Unit), and its variants. These functions have different properties and are chosen based on the specific requirements of the task and the network architecture. For example, ReLU is widely used in deep neural networks due to its ability to alleviate the vanishing gradient problem and promote sparsity in the activations. By applying

activation functions, the neural network can learn capture non-linear relationships in the data, enabling it to model complex patterns and make accurate predictions or decisions.

(54) The encoder **310** breaks the decompressed outputs passed by the decompressed output organizer **190** down into smaller representations of the original data sets. Following the encoder the data may pass through a transformer **330**. A transformer is a type of neural network architecture that may rely on a self-attention mechanism which allows the model to weigh the importance of different parts of the input sequence when processing each element. This enables the transformer to capture dependencies and relationships between elements in the sequence efficiently. After being processed by a transformer **330**, the data sets may be further processed by a decoder **320** which restores the smaller representations back into the original decompressed data sets. The decoder **320** may have a similar composition as the encoder **310**, but reversed, to undo the operations performed on the data sets by the encoder **310**. The transformer **330** may identify important aspects in each group of decompressed data passed through the correlation network which allows the decoder **320** to rebuild a more complete version of the original decompressed data sets. The decoder **320** may output an N number of restored data sets **340** which correspond to the N number of correlated data sets **300** originally passed through the correlation network **170**.

(55) FIG. 4 is a block diagram illustrating an exemplary aspect of a platform for a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, an autoencoder training system **270**. According to the embodiment, the autoencoder training system **270** may comprise a model training stage comprising a data preprocessor **402**, one or more machine and/or deep learning algorithms **403**, training output **404**, and a parametric optimizer **405**, and a model deployment stage comprising a deployed and fully trained model **410** configured to perform tasks described herein such as transcription, summarization, agent coaching, and agent guidance. Autoencoder training system **270** may be used to train and deploy a multi-layer autoencoder network in order to support the services provided by the compression and restoration system.

(56) At the model training stage, a plurality of training data **401** may be received at the autoencoder training system **270**. In some embodiments, the plurality of training data may be obtained from one or more storage systems and/or directly from various information sources. In a use case directed to hyperspectral images, a plurality of training data may be sourced from data collectors including but not limited to satellites, airborne sensors, unmanned aerial vehicles, ground-based sensors, and medical devices. Hyperspectral data refers to data that includes wide ranges of the electromagnetic spectrum. It could include information in ranges including but not limited to the visible spectrum and the infrared spectrum. Data preprocessor **402** may receive the input data (e.g., hyperspectral data) and perform various data preprocessing tasks on the input data to format the data for further processing. For example, data preprocessing can include, but is not limited to, tasks related to data cleansing, data deduplication, data normalization, data transformation, handling missing values, feature extraction and selection, mismatch handling, and/or the like. Data preprocessor **402** may also be configured to create training dataset, a validation dataset, and a test set from the plurality of input data **401**. For example, a training dataset may comprise 80% of the preprocessed input data, the validation set 10%, and the test dataset may comprise the remaining 10% of the data. The preprocessed training dataset may be fed as input into one or more machine and/or deep learning algorithms **403** to train a predictive model for object monitoring and detection.

(57) During model training, training output **404** is produced and used to measure the quality and efficiency of the compressed outputs. During this process a parametric optimizer **405** may be used to perform algorithmic tuning between model training iterations. Model parameters and hyperparameters can include, but are not limited to, bias, train-test split ratio, learning rate in optimization algorithms (e.g., gradient descent), choice of optimization algorithm (e.g., gradient descent, stochastic gradient descent, of Adam optimizer, etc.), choice of activation function in a neural network layer (e.g., Sigmoid, ReLU, Tan h, etc.), the choice of cost or loss function the model will use, number of hidden layers in a neural network, number of activation unites in each

layer, the drop-out rate in a neural network, number of iterations (epochs) in a training the model, number of clusters in a clustering task, kernel or filter size in convolutional layers, pooling size, batch size, the coefficients (or weights) of linear or logistic regression models, cluster centroids, and/or the like. Parameters and hyperparameters may be tuned and then applied to the next round of model training. In this way, the training stage provides a machine learning training loop.

(58) In some implementations, various accuracy metrics may be used by the autoencoder training system **270** to evaluate a model's performance. Metrics can include, but are not limited to, compression ratio, the amount of data lost, the size of the compressed file, and the speed at which data is compressed, to name a few. In one embodiment, the system may utilize a loss function **407** to measure the system's performance. The loss function **407** compares the training outputs with an expected output and determined how the algorithm needs to be changed in order to improve the quality of the model output. During the training stage, all outputs may be passed through the loss function **407** on a continuous loop until the algorithms **403** are in a position where they can effectively be incorporated into a deployed model **415**.

(59) The test dataset can be used to test the accuracy of the model outputs. If the training model is compressing or decompressing data to the user's preferred standards, then it can be moved to the model deployment stage as a fully trained and deployed model **410** in a production environment compressing or decompressing live input data **411** (e.g., hyperspectral data). Further, model compressions or decompressions made by deployed model can be used as feedback and applied to model training in the training stage, wherein the model is continuously learning over time using both training data and live data and predictions.

(60) A model and training database **406** is present and configured to store training/test datasets and developed models. Database **406** may also store previous versions of models. According to some embodiments, the one or more machine and/or deep learning models may comprise any suitable algorithm known to those with skill in the art including, but not limited to: LLMs, generative transformers, transformers, supervised learning algorithms such as: regression (e.g., linear, polynomial, logistic, etc.), decision tree, random forest, k-nearest neighbor, support vector machines, Naïve-Bayes algorithm; unsupervised learning algorithms such as clustering algorithms, hidden Markov models, singular value decomposition, and/or the like. Alternatively, or additionally, algorithms **403** may comprise a deep learning algorithm such as neural networks (e.g., recurrent, convolutional, long short-term memory networks, etc.). In some implementations, the autoencoder training system **270** automatically generates standardized model scorecards for each model produced to provide rapid insights into the model and training data, maintain model provenance, and track performance over time. These model scorecards provide insights into model framework(s) used, training data, training data specifications such as chip size, stride, data splits, baseline hyperparameters, and other factors. Model scorecards may be stored in database(s) **406**.

(61) FIG. 5 is a block diagram illustrating an exemplary aspect of a subsystem of the system for compressing and restoring data using multi-level autoencoders and correlation networks, a correlation network training system **370**. According to the embodiment, correlation network training system **370** may comprise a model training stage comprising a data preprocessor **502**, one or more machine and/or deep learning algorithms **503**, training output **504**, and a parametric optimizer **505**, and a model deployment stage comprising a deployed and fully trained model **510** configured to perform tasks described herein such determining correlations between compressed data sets. The correlation network training system **370** may be used to train and deploy the correlation network **300** in order to support the services provided by the compression and decompression system.

(62) At the model training stage, a plurality of training data **501** may be received by the correlation network training system **370**. In some embodiments, the plurality of training data may be obtained from one or more storage systems and/or directly from the compression network **131**. In some embodiments, the correlation network training system may obtain data sets from a vector grouping

system. In a use case directed to hyperspectral data sets, a plurality of decompressed training data may be sourced from a hyperspectral data compression system. Data preprocessor **502** may receive the input data (e.g., decompressed hyperspectral data) and perform various data preprocessing tasks on the input data to format the data for further processing. For example, data preprocessing can include, but is not limited to, tasks related to data cleansing, data deduplication, data normalization, data transformation, handling missing values, feature extraction and selection, mismatch handling, and/or the like. Data preprocessor **502** may also be configured to create training dataset, a validation dataset, and a test set from the plurality of input data **501**. For example, a training dataset may comprise 80% of the preprocessed input data, the validation set 10%, and the test dataset may comprise the remaining 10% of the data. The preprocessed training dataset may be fed as input into one or more machine and/or deep learning algorithms **503** to train a predictive model for object monitoring and detection.

(63) During model training, training output **504** is produced and used to measure the accuracy and usefulness of the predictive outputs. During this process a parametric optimizer **505** may be used to perform algorithmic tuning between model training iterations. Model parameters and hyperparameters can include, but are not limited to, bias, train-test split ratio, learning rate in optimization algorithms (e.g., gradient descent), choice of optimization algorithm (e.g., gradient descent, stochastic gradient descent, of Adam optimizer, etc.), choice of activation function in a neural network layer (e.g., Sigmoid, ReLU, Tan h, etc.), the choice of cost or loss function the model will use, number of hidden layers in a neural network, number of activation unites in each layer, the drop-out rate in a neural network, number of iterations (epochs) in a training the model, number of clusters in a clustering task, kernel or filter size in convolutional layers, pooling size, batch size, the coefficients (or weights) of linear or logistic regression models, cluster centroids, and/or the like. Parameters and hyperparameters may be tuned and then applied to the next round of model training. In this way, the training stage provides a machine learning training loop.

(64) In some implementations, various accuracy metrics may be used by a machine learning engine to evaluate a model's performance. Metrics can include, but are not limited to, word error rate (WER), word information loss, speaker identification accuracy (e.g., single stream with multiple speakers), inverse text normalization and normalization error rate, punctuation accuracy, timestamp accuracy, latency, resource consumption, custom vocabulary, sentence-level sentiment analysis, multiple languages supported, cost-to-performance tradeoff, and personal identifying information/payment card industry redaction, to name a few. In one embodiment, the system may utilize a loss function **507** to measure the system's performance. The loss function **507** compares the training outputs with an expected output and determined how the algorithm needs to be changed in order to improve the quality of the model output. During the training stage, all outputs may be passed through the loss function **507** on a continuous loop until the algorithms **503** are in a position where they can effectively be incorporated into a deployed model **515**.

(65) The test dataset can be used to test the accuracy of the model outputs. If the training model is establishing correlations that satisfy a certain criterion such as but not limited to quality of the correlations and amount of restored lost data, then it can be moved to the model deployment stage as a fully trained and deployed model **510** in a production environment making predictions based on live input data **511** (e.g., compressed hyperspectral data). Further, model correlations and restorations made by deployed model can be used as feedback and applied to model training in the training stage, wherein the model is continuously learning over time using both training data and live data and predictions. A model and training database **506** is present and configured to store training/test datasets and developed models. Database **506** may also store previous versions of models.

(66) According to some embodiments, the one or more machine and/or deep learning models may comprise any suitable algorithm known to those with skill in the art including, but not limited to: LLMs, generative transformers, transformers, supervised learning algorithms such as: regression

(e.g., linear, polynomial, logistic, etc.), decision tree, random forest, k-nearest neighbor, support vector machines, Naïve-Bayes algorithm; unsupervised learning algorithms such as clustering algorithms, hidden Markov models, singular value decomposition, and/or the like. Alternatively, or additionally, algorithms **503** may comprise a deep learning algorithm such as neural networks (e.g., recurrent, convolutional, long short-term memory networks, etc.).

(67) In some implementations, the correlation network training system **270** automatically generates standardized model scorecards for each model produced to provide rapid insights into the model and training data, maintain model provenance, and track performance over time. These model scorecards provide insights into model framework(s) used, training data, training data specifications such as chip size, stride, data splits, baseline hyperparameters, and other factors. Model scorecards may be stored in database(s) **506**.

(68) FIG. **13** is a block diagram **1300** illustrating an exemplary system architecture for compressing and restoring multi-modal data using a variational autoencoder. In one or more embodiments, compressing and restoring multi-modal data using a variational autoencoder enables homomorphic data compression across diverse data types. The multi-modal variational autoencoder can be configured to learn a smooth, continuous unified latent space representation of the multi-modal input data. Thus, in embodiments, small perturbations in the unified latent space correspond to significant variations in the input data across modalities. In embodiments, linear operations in the unified latent space, such as addition, subtraction, and scalar multiplication correspond to meaningful operations in the multi-modal input data.

(69) In general, data compression has advantages for computer systems in terms of resource usage and scalability. Data compression techniques can significantly reduce the storage space required for data while still maintaining its integrity and utility. This is particularly valuable in applications dealing with large volumes of data, such as cloud computing and big data analytics. Moreover, using data compression can lead to more efficient transmission over networks, reducing bandwidth requirements.

(70) Homomorphic data compression offers several additional benefits over conventional data compression techniques, primarily in scenarios where preserving privacy and reducing computational overhead are important. One of the primary advantages of homomorphic data compression is the ability to compress data while keeping it encrypted. This is especially important in fields such as healthcare and finance, where sensitive data needs to be analyzed without compromising privacy. Homomorphic compression enables secure outsourcing of computations to third-party service providers while preserving data privacy. This is advantageous in scenarios where organizations need to leverage external resources for data processing without exposing sensitive information. Homomorphic compression techniques can facilitate data analysis on encrypted data without the need for decryption, enabling organizations to derive valuable insights while maintaining data confidentiality. In industries with stringent data privacy regulations, such as healthcare (Health Insurance Portability and Accountability Act of the United States (HIPAA)) and finance (General Data Protection Regulation of the European Union) (GDPR)), homomorphic compression can help organizations comply with regulatory requirements by ensuring the privacy and security of sensitive information.

(71) In one embodiment, a system for compressing and restoring multi-modal data using a variational autoencoder comprises a plurality of multi-modal data inputs **1302**, a multi-modal data preprocessor **1310**, a data normalizer **1320**, a multi-modal variational autoencoder **1330** which further comprises an encoder network **1331** and a decoder network **1332**. Configured and disposed between the encoder network **1331** and the decoder network **1332** is a unified latent space **1340**. The unified latent space **1340** represents the low-dimensional continuous space where the autoencoder maps the multi-modal input data. The encoder network **1331** within the multi-modal variational autoencoder **1330** receives multi-modal data input **1302**. Data input **1302** can include text, images, audio, video, hyperspectral data, time-series data, sensor data, binary data, and/or

other types of data. The encoder network **1331** maps the received multi-modal data input **1302** to a probability distribution in the unified latent space **1340**. In embodiments, the encoder network **1331** outputs parameters of a probability distribution. In embodiments, the parameters can include mean and variance. In embodiments, the probability distribution includes a Gaussian (normal) distribution. In one or more embodiments, the probability distribution can include a Bernoulli distribution, categorical distribution, and/or a multinomial distribution, instead of, or in addition to, a Gaussian distribution. In embodiments, programming instructions executing on a processor cause the variational autoencoder to perform partially homomorphic compression and/or partially homomorphic encryption. In embodiments, the variational autoencoder comprises a latent space, and wherein the plurality of programming instructions further includes instructions that, when executing on the processor, cause the variational autoencoder to perform one or more linear operations in the latent space. In embodiments, programming instructions further includes instructions that, when executing on the processor, cause the variational autoencoder to perform the linear operations that include at least one of addition, subtraction, and scalar multiplication.

(72) The decoder network **1332** maps the unified latent space **1340** to a decompressed multi-modal output **1370**. The decompressed multi-modal output can be input to a decompressed multi-modal output organizer **1390**, which in turn outputs a reconstructed multi-modal output **1380**. The multi-modal data inputs **1302** can be representations of raw data from various sources, such as sensors, cameras, microphones, or databases. The raw data can be in different formats, including but not limited to images, videos, audio, time-series, or structured data. The plurality of multi-modal data inputs **1302** may be transferred to the multi-modal data preprocessor **1310** for further processing. The multi-modal data preprocessor **1310** applies various preprocessing techniques to the raw data received from the data input **1302**. These techniques may include data cleaning, noise reduction, artifact removal, format conversion, and synchronization of different data modalities. The multi-modal data preprocessor **1310** ensures that the data from all modalities is in a suitable format and quality for subsequent stages of the system.

(73) The preprocessed data may then be passed to the data normalizer **1320**. The data normalizer **1320** scales and normalizes the data to a consistent range, typically between 0 and 1. Normalization can help to improve the training stability and convergence of the multi-modal variational autoencoder **1330**. The normalized data is fed into the multi-modal variational autoencoder **1330**, which includes both the encoder network **1331**, latent space **1340**, and the decoder network **1332**. The encoder network **1331** is responsible for encoding the input data into a lower-dimensional latent space representation. The decoder network **1332** can include multiple layers of encoders that progressively reduce the dimensionality of the data while capturing the most important features and patterns.

(74) The latent space **1340** has a significantly reduced size compared to the original input data, enabling efficient storage and transmission. The latent space **1340** represents a compressed output that may be stored in a storage system. A storage system may include any suitable storage medium, such as a database, file system, or cloud storage. Storage systems allow for the efficient management and retrieval of the compressed data represented by latent space **1340** as needed.

(75) In embodiments, one or more operations may be performed on the compressed data of latent space **1340**. The operations can include homomorphic operations. The homomorphic operations can include linear operations. The linear operations can include addition, subtraction, scalar multiplication, and/or other suitable operations.

(76) When the compressed data needs to be restored or reconstructed, it may be retrieved from the storage system and passed to the decoder network **1332**. Additionally, the compressed data may be directly passed directly to the decoder network **1332**. The decoder network **1332** is responsible for decoding the compressed latent representation back into the original data space by outputting a decompressed output **1370**. The decoder network **1332** can include multiple layers of decoders that progressively increase the dimensionality of the data, reconstructing the original input.

(77) The decompressed multi-modal output **1370** from the decoder network **1332** may have some loss of information compared to the original multi-modal input data due to the compression process. In one or more embodiments, to further enhance the quality of the decompressed output, the system may optionally incorporate a multi-modal correlation network **1360**. The multi-modal correlation network **1360** leverages the correlations and patterns between different compressed inputs, both within and across modalities, to restore the decompressed output more accurately. It learns to capture the relationships and dependencies within and between different data types, allowing for better reconstruction and restoration of the original multi-modal information. The multi-modal correlation network **1360** takes the decompressed multi-modal outputs **1370** as inputs. It analyzes the correlations and similarities between the data samples across modalities and uses this information to refine and enhance the decompressed output. The refined decompressed output from the multi-modal correlation network **1360** is a reconstructed multi-modal output **1380** of the system. The reconstructed multi-modal output **1380** closely resembles the original multi-modal input data, with minimal loss of information and improved quality compared to the output from the decoder network **1332** alone.

(78) In one or more embodiments, the correlation network **1360** may receive inputs from a decompressed output organizer **1390** which operates on the decompressed outputs **1370** obtained from the decoder network **1332**. The decompressed output organizer **1390** may organize the decompressed outputs **1370** into groups based on their correlations and similarities. By grouping decompressed outputs **1370** based on similarities, the correlation network **1360** will more easily be able to identify correlations between decompressed outputs **1370**. Thus, embodiments can further include a correlation network, and wherein the output of the decompressed output organizer is input to the correlation network, and wherein the output of the correlation network provides the reconstructed output.

(79) In one or more embodiments, the variational autoencoder **1330** is trained via training data **1372**. The training data **1372** can be part of a supervised or semi-supervised learning process. In embodiments, the variational autoencoder **1330** can be optimized with an ADAM (Adaptive Moment Estimation) optimizer **1374**. In one or more embodiments, the encoder network **1331** and decoder network **1332** are trained to learn meaningful representations of the data input **1302** and generate realistic output data for reconstructed output **1380**. In embodiments, the ADAM optimizer **1374** is used to adjust the weights of the encoder network **1331** and/or decoder network **1332** based on factors such as the reconstruction error and/or the KL divergence between the learned latent distribution and a predefined prior distribution, such as a Gaussian distribution. In one or more embodiments, the KL (Kullback-Leibler) divergence is used as part of the objective function in probabilistic models to encourage learned distributions to match target distributions, regularize the model, and/or learn meaningful representations to enable features such as homomorphic data compression. Thus, embodiments can include training the variational autoencoder, wherein the training includes using an Adam optimizer.

(80) FIG. **14** is a block diagram **1400** of a multi-modal variational autoencoder used in an embodiment. The general architecture of a multi-modal autoencoder includes a multi-modal encoder, multi-modal decoder, and unified latent space. The unified latent space may be implemented using a bottleneck layer. The encoder section includes multiple input layers that accept different types of input data. Hidden layers progressively reduce the dimensionality of the multi-modal input, capturing important features and patterns across all modalities. These layers comprise the multi-modal encoder. The bottleneck layer (unified latent space) is the final hidden layer, where the dimensionality is significantly reduced. This layer represents the compressed encoding of the multi-modal input data in a shared representation.

(81) The multi-modal input data **1402** can include various types of uncompressed data. In one or more embodiments, the uncompressed data can include hyperspectral data, image data, audio data, text data, time-series data, and sensor data. The sensor data can originate from a wide variety of

sensors, including IoT sensors. The multi-modal input data **1402** is provided to a multi-modal probabilistic encoder **1404**. The multi-modal probabilistic encoder **1404** is configured to process different data types using modality-specific layers and then integrate this information using shared layers. It can be configured to output a deterministic encoding of the multi-modal input data **1402** as well as a probability distribution to generate a unified latent space **1406**. The unified latent space **1406** may have homomorphic properties, which can enable homomorphic data compression and/or partial homomorphic data compression across all modalities. The information from the unified latent space **1406** can be input to a multi-modal probabilistic decoder **1408**. In one or more embodiments, the multi-modal probabilistic decoder **1408** outputs the parameters of a probability distribution for each data modality. These distributions are then sampled to obtain a reconstructed multi-modal output **1410**. The multi-modal variational autoencoder may utilize cross-modal attention mechanisms to capture relationships between different data types during both encoding and decoding processes. In one or more embodiments, the multi-modal variational autoencoder may have a layer configuration similar to that shown in FIG. 2, where the innermost layer that provides compressed output (**140** of FIG. 2) serves as a bottleneck layer, but with additional modality-specific layers and cross-modal attention mechanisms to handle diverse data types.

(82) FIG. 15 is a flow diagram **1500** illustrating an exemplary method for compressing multi-modal data using a system for compressing and restoring data using multi-modal variational autoencoders, according to an embodiment. In a first step **1510**, a plurality of multi-modal data sets is collected from a plurality of data sources. These data sources can include various sensors, devices, databases, or any other systems that generate or store diverse types of data. The data sets are inherently multi-modal, meaning they comprise different data types and formats. For example, the multi-modal data sets can include combinations of images, videos, audio recordings, time-series data, numerical data, textual data, and sensor readings. The collection process involves acquiring these diverse data sets from their respective sources and bringing them into a centralized system for further processing.

(83) In step **1520**, the collected multi-modal data sets are preprocessed using a multi-modal data preprocessor. The multi-modal data preprocessor is responsible for steps including one or more of cleaning, transforming, normalizing, aligning, and preparing the diverse data sets for subsequent analysis and compression. Preprocessing tasks may include, but are not limited to, data cleansing, data integration, data transformation, feature extraction, synchronization and alignment of different data modalities. Data cleansing involves removing or correcting any erroneous, missing, or inconsistent data points across all modalities. Data integration combines data from multiple sources and modalities into a unified format. Data transformation converts the various data types into suitable representations for further processing, such as scaling, normalization, or encoding categorical variables. Feature extraction identifies and selects relevant features or attributes from the multi-modal data sets that are most informative for the given task. The preprocessor also ensures that data from different modalities are properly aligned and synchronized for joint processing.

(84) In step **1530**, the preprocessed multi-modal data sets are compressed into a compressed output using a multi-modal variational autoencoder. The multi-modal variational autoencoder provides a deep learning model designed to learn compact and meaningful representations of the diverse input data in a unified latent space. As previously described and shown in at least FIG. 13 and FIG. 14, the multi-modal variational autoencoder of disclosed embodiments includes a multi-modal encoder network and a multi-modal decoder network. The multi-modal encoder network takes the normalized multi-modal data sets as input and progressively compresses them through a series of modality-specific and shared layers. These may include, but are not limited to, convolutional layers for image data, recurrent layers for sequential data, and fully connected layers for numerical data. The compressed representation is obtained at the bottleneck layer of the encoder network, which represents a unified latent space with significantly reduced dimensionality compared to the original multi-modal data. The multi-modal variational autoencoder may utilize cross-modal attention

mechanisms to capture relationships between different data types. During the compression process, the multi-modal variational autoencoder learns to capture the essential features and patterns present in the diverse data sets while discarding redundant or irrelevant information. It aims to minimize the reconstruction error between the original multi-modal data and the reconstructed data obtained from the compressed representation, accounting for the characteristics of each data modality. In step **1540**, the compressed output generated by the multi-modal variational autoencoder is outputted and/or stored for future processing. This compressed output represents the compact and informative representation of the original multi-modal data sets in a unified latent space. The compressed output can support homomorphic operations, enabling manipulation on encrypted multi-modal data. It significantly reduces the storage and transmission requirements compared to the original diverse data sets, making it more efficient for downstream tasks while preserving the relationships between different data modalities.

(85) FIG. **16** is a flow diagram **1600** illustrating an exemplary method for decompressing multi-modal data input using a system for compressing and restoring data using multi-modal variational autoencoders, according to an embodiment. In a first step, **1610**, a plurality of compressed multi-modal data sets is accessed from a unified latent space. Step **1620** includes decompressing the plurality of compressed multi-modal data sets using a multi-modal variational autoencoder, such as previously described and shown in at least FIG. **13** and/or FIG. **14**. The multi-modal decoder network is responsible for mapping the unified latent space vectors back to the original multi-modal data space. The decoder network may include modality-specific techniques such as transposed convolutions for image data, recurrent layers for sequential data, and fully connected layers for numerical data, as well as shared layers and cross-modal attention mechanisms. In one or more embodiments, homomorphic operations may be performed on the compressed multi-modal data sets in the unified latent space prior to decompressing.

(86) In a step **1630**, the compressed multi-modal data sets are reconstructed using the multi-modal decoder (e.g., **1332** of FIG. **13**) of the variational autoencoder. In a step **1640**, the restored, decompressed multi-modal data set is outputted. The restored multi-modal data set represents the reconstructed version of the original diverse data types, which includes recovered information lost during the compression process across all modalities. The outputted multi-modal data may have had homomorphic operations performed on the unified latent space compressed version. This feature enables meaningful operations on multi-modal data while in its compressed (and/or encrypted) form which still apply when the data is decompressed, preserving relationships and transformations across different data types.

(87) FIG. **17** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using multi-level autoencoders and correlation networks on a plurality of edge devices. The system comprises two main segments: an edge device component that handles initial data processing and compression, and a cloud infrastructure component that manages decompression and reconstruction.

(88) The process begins with multi-modal data input **1302** which represents raw data streams from various sources such as sensors, cameras, or databases. The system supports deployment across multiple edge devices, each handling different aspects of data collection and processing. For example, in a smart manufacturing facility, edge devices might include industrial cameras performing quality control, vibration sensors monitoring equipment health, thermal sensors tracking operating temperatures, and audio sensors detecting machine anomalies. Each edge device operates independently while maintaining synchronization with the broader system.

(89) Within each edge device **1700**, a multi-modal data preprocessor **1310** applies various preprocessing techniques optimized for that device's specific deployment. Consider an autonomous vehicle system where multiple edge devices work in concert. The forward-facing cameras preprocess visual data for obstacle detection, while LIDAR sensors handle point cloud data for distance measurements. Simultaneously, microphone arrays preprocess audio for emergency

vehicle detection while internal sensors manage vehicle telemetry data. These diverse data streams are processed in parallel across the distributed edge devices, each optimized for its specific data type while maintaining synchronization with the broader system.

(90) The preprocessor ensures that multi-modal data is formatted appropriately while minimizing computational overhead on each edge device. In a smart building deployment, different edge devices might handle various aspects of the building's operation. Some edge devices focus on environmental monitoring, processing temperature, humidity, and air quality sensor data. Others manage security systems, handling camera feeds and access control data. Additional devices might process energy management data, including power consumption and occupancy information, while network monitoring devices handle connectivity data. Each edge device applies specific preprocessing optimizations based on its designated role and the types of data it handles.

(91) This distributed approach allows for scalable deployment where additional edge devices can be seamlessly integrated into the system as needed. For instance, in a smart city deployment, edge devices might be added incrementally as new sensors and monitoring systems are installed. A traffic monitoring system might start with basic camera-equipped edge devices, then expand to include devices processing data from traffic flow sensors, environmental monitors, and emergency response systems. Each new edge device integrates into the existing network, contributing to the overall multi-modal data processing pipeline while maintaining efficient resource utilization and data synchronization across the entire system.

(92) A data normalizer **1320** scales and normalizes the preprocessed data to ensure consistent ranges across all modalities. This normalization step enables stable processing in the subsequent compression stages and helps maintain data quality while operating within edge device constraints. For edge devices processing multiple data types, normalization provides a foundation for effective compression by bringing diverse data formats into aligned numerical ranges. In a manufacturing setting, for example, the normalizer processes temperature readings that might range from 0 to 1000 degrees, vibration measurements spanning microscopic distances, and audio signals with varying amplitudes, mapping each into standardized ranges that the compression system can efficiently handle.

(93) Data normalizer **1320** is capable of adapting its scaling parameters based on the specific requirements of each edge device and its data types. In autonomous vehicle systems, the normalizer handles visual data from cameras, ranging from bright daylight to near-darkness conditions, alongside LIDAR measurements that might span several hundred meters, and acceleration data that varies by fractions of a g-force. Each data stream is scaled appropriately while preserving the relative relationships between measurements, ensuring that no information is lost during the normalization process.

(94) In one embodiment, the normalization process may also consider the computational limitations of edge devices. Rather than performing complex statistical transformations that might overtax limited processing resources, the normalizer employs efficient scaling methods optimized for edge deployment. This approach allows edge devices to maintain high throughput while ensuring data quality. For instance, in IoT sensor networks, where devices might operate on battery power or limited computing resources, the normalizer balances the need for consistent data scaling with power and processing constraints.

(95) This standardization of data ranges across modalities serves as a prerequisite for the subsequent compression stages, enabling the edge encoder to operate efficiently on normalized data regardless of its original scale or format. The process maintains the statistical relationships within each data type while creating a unified representation that can be effectively compressed and later reconstructed by the cloud infrastructure.

(96) An edge encoder network **1710** compresses the normalized data using an optimized implementation of the encoder portion of the variational autoencoder, specifically designed for efficient operation on edge devices. The entire architecture is specifically optimized for deployment

on resource-constrained edge devices through several key design considerations. The encoder network employs a lightweight implementation that minimizes computational overhead while maintaining compression efficiency, utilizing optimized layer configurations and reduced parameter counts compared to traditional variational autoencoders. The system dynamically adjusts its processing based on each edge device's available resources, scaling the complexity of operations to match local computational capabilities, memory constraints, and power limitations. For instance, the preprocessing and normalization stages employ efficient algorithms that minimize memory usage, while the edge latent space is designed to achieve high compression ratios with minimal computational cost. The architecture supports local caching and batched processing to manage resource utilization effectively, allowing edge devices to compress and process data even with limited memory or intermittent connectivity. This resource-aware design extends to the homomorphic operations, which may be implemented using optimized algorithms suitable for edge computing constraints while still maintaining their privacy-preserving properties. In one embodiment the system may employ adaptive sampling and processing rates based on device capabilities and current load, ensuring stable operation across a wide range of edge devices from simple IoT sensors to more powerful edge computing nodes.

(97) Compressed representations created by edge encoder network **1710** is maintained in an edge latent space **1720**, which provides a unified representation of all data modalities while enabling homomorphic operations. This unified latent space serves as a mathematical bridge between different data types, allowing for operations to be performed on the compressed data without decompression. For instance, in a smart building system, the latent space might encode relationships between occupancy patterns, energy usage, and environmental conditions in a way that allows for direct analysis and optimization in the compressed domain. The homomorphic properties of this latent space enable privacy-preserving computations, allowing edge devices to perform meaningful operations on sensitive data while maintaining encryption.

(98) Compressed data **140** from the edge latent space is then transmitted to the cloud infrastructure **1730**, implementing an efficient data transfer protocol that minimizes bandwidth usage while ensuring data integrity. This transmission process can adapt to varying network conditions, potentially batching compressed representations during periods of limited connectivity or prioritizing certain data types based on application requirements. For example, in a medical monitoring system, critical patient data might be prioritized for immediate transmission while less time-sensitive data is batched for later transfer.

(99) Once received by the cloud infrastructure, the cloud decoder network **1740** begins the reconstruction process, employing decompression techniques that leverage the full computational resources available in the cloud environment. This network maps the compressed representations back to their original data spaces through a series of expanding layers that progressively rebuild the multi-modal data. The decoder **1740** maintains the complex relationships between different modalities established during compression, ensuring that reconstructed data maintains its inter-modal correlations. For instance, in an autonomous vehicle system, the decoder might reconstruct visual scene data alongside LIDAR measurements and vehicle telemetry in a way that preserves the spatial and temporal relationships critical for accurate environment understanding and decision-making.

(100) A cloud decompressed output organizer **1750** arranges the decompressed outputs based on their cross-modal correlations and similarities. This organization facilitates the correlation network's **1760** ability to identify and leverage relationships between different data modalities, allowing it to recover information that may have been lost during compression. Correlation network **1760** analyzes the organized outputs to identify patterns and dependencies both within and across modalities. By leveraging these relationships, it enhances the quality of the reconstructed data. The system ultimately produces a reconstructed output **1770** that closely resembles the original multi-modal input while maintaining the essential relationships between different data

types.

(101) This distributed architecture enables efficient processing of multi-modal data by optimizing resource usage between edge devices and cloud infrastructure. For example, in a smart manufacturing environment, edge devices could compress sensor data, machine vision inputs, and audio measurements from equipment, while the cloud infrastructure handles the complex task of reconstructing and analyzing this multi-modal data to detect patterns or anomalies.

(102) FIG. **18** is a block diagram illustrating an exemplary system architecture for compressing and restoring data using variational autoencoders, according to an embodiment. In one embodiment, the system may incorporate training components that optimize the system's performance across both edge devices and cloud infrastructure. Multi-modal training data **1372** serves as the foundation for developing compression capabilities, comprising diverse examples that span multiple data modalities. This training data encompasses combinations such as sensor readings, visual inputs, audio signals, and textual data, enabling the system to learn both modality-specific features and cross-modal relationships.

(103) An ADAM optimizer **1374** represents an optimization approach well-suited for training deep learning models in resource-constrained edge environments. By utilizing estimates of first and second moments of the gradients, it adaptively adjusts learning rates for different parameters, ensuring stable and efficient training despite the limitations of edge hardware. The optimizer interacts directly with the edge encoder network **1710** to refine its compression capabilities while preserving the network's ability to generate meaningful representations in the edge latent space **1720**.

(104) During the training process, the system learns to minimize reconstruction loss between the original multi-modal input data and the decompressed output while accounting for reconstruction quality across all data modalities. ADAM optimizer **1374** helps achieve this by dynamically adjusting the learning process based on the computed gradients, allowing the system to learn efficient compression strategies that work within the computational constraints of edge devices.

(105) The training approach incorporates both supervised and semi-supervised learning processes, allowing the system to leverage both labeled and unlabeled data effectively. ADAM optimizer **1374** uses factors such as reconstruction error and KL divergence between learned latent distributions and predefined prior distributions to adjust the weights of the edge encoder network. This optimization process helps ensure that the compressed representations maintain meaningful relationships between different data modalities while achieving significant data reduction.

(106) In practical applications, this training architecture enables continuous improvement and adaptation to specific deployment scenarios. For instance, in a smart city environment, the training data might include traffic camera feeds, environmental sensor readings, and urban noise measurements. ADAM optimizer **1374** ensures the edge encoder network learns to effectively compress this diverse data while maintaining the relationships necessary for meaningful analysis in the cloud infrastructure. The system can be periodically retrained with new data to adapt to changing patterns and requirements, with the optimizer ensuring efficient updates to the model parameters.

(107) The combination of carefully curated multi-modal training data and adaptive optimization enables the system to achieve high compression ratios while preserving essential information across all data modalities. This approach is particularly valuable in scenarios where edge devices must process diverse data types under resource constraints while maintaining the ability to reconstruct meaningful insights in the cloud infrastructure.

(108) Detailed Description of Exemplary Aspects

(109) FIG. **6** is a flow diagram illustrating an exemplary method for compressing a data input using a system for compressing and restoring data using multi-level autoencoders and correlation networks. In a first step **600**, a plurality of data sets is collected from a plurality of data sources. These data sources can include various sensors, devices, databases, or any other systems that

generate or store data. The data sets may be heterogeneous in nature, meaning they can have different formats, structures, or modalities. For example, the data sets can include images, videos, audio recordings, time-series data, numerical data, or textual data. The collection process involves acquiring the data sets from their respective sources and bringing them into a centralized system for further processing.

(110) In a step **610**, the collected data sets are preprocessed using a data preprocessor. The data preprocessor may be responsible for cleaning, transforming, and preparing the data sets for subsequent analysis and compression. Preprocessing tasks may include but are not limited to data cleansing, data integration, data transformation, and feature extraction. Data cleansing involves removing or correcting any erroneous, missing, or inconsistent data points. Data integration combines data from multiple sources into a unified format. Data transformation converts the data into a suitable representation for further processing, such as scaling, normalization, or encoding categorical variables. Feature extraction identifies and selects relevant features or attributes from the data sets that are most informative for the given task.

(111) A step **620** involves normalizing the preprocessed data sets using a data normalizer. Normalization is a step that brings the data into a common scale and range. It helps to remove any biases or inconsistencies that may exist due to different units or scales of measurement. The data normalizer applies various normalization techniques, such as min-max scaling, z-score normalization, or unit vector normalization, depending on the nature of the data and the requirements of the subsequent compression step. Normalization ensures that all the data sets have a consistent representation and can be compared and processed effectively.

(112) In a step **630**, the normalized data sets are compressed into a compressed output using a multi-layer autoencoder network. The multi-layer autoencoder network is a deep learning model designed to learn compact and meaningful representations of the input data. It consists of an encoder network and a decoder network. The encoder network takes the normalized data sets as input and progressively compresses them through a series of layers, such as but not limited to convolutional layers, pooling layers, and fully connected layers. The compressed representation is obtained at the bottleneck layer of the encoder network, which has a significantly reduced dimensionality compared to the original data. The multi-layer autoencoder network may utilize a plurality of encoder networks to achieve optimal compression performance. These encoder networks can include different architectures, loss functions, or optimization techniques. The choice of compression technique depends on the specific characteristics and requirements of the data sets being compressed. During the compression process, the multi-layer autoencoder network learns to capture the essential features and patterns present in the data sets while discarding redundant or irrelevant information. It aims to minimize the reconstruction error between the original data and the reconstructed data obtained from the compressed representation. In step **640**, the compressed output generated by the multi-layer autoencoder network is either outputted or stored for future processing. The compressed output represents the compact and informative representation of the original data sets. It can be transmitted, stored, or further analyzed depending on the specific application or use case. The compressed output significantly reduces the storage and transmission requirements compared to the original data sets, making it more efficient for downstream tasks.

(113) FIG. 7 is a flow diagram illustrating an exemplary method for decompressing a compressed data input using a system for compressing and restoring data using multi-level autoencoders and correlation networks. In a first step, **700**, access a plurality of compressed data sets. In a step **710**, decompress the plurality of compressed data sets using a multi-layer autoencoder's decoder network. The decoder network is responsible for mapping the latent space vectors back to the original data space. The decoder network may include techniques such as transposed convolutions, upsampling layers, or generative models, depending on the specific requirements of the data and the compression method used.

(114) Step **720** includes leveraging the similarities between decompressed outputs using a

correlation network, which may exploit shared information and patterns to achieve a better reconstruction. The correlation network is a deep learning model specifically designed to exploit the shared information and patterns among the compressed data sets. It takes the organized decompressed data sets as input and learns to capture the correlations and dependencies between them. The correlation network may consist of multiple layers, such as convolutional layers, recurrent layers, or attention mechanisms, which enable it to effectively model the relationships and similarities among the compressed data sets.

(115) In a step **730**, the compressed data sets are reconstructed using the correlation network. The reconstruction process in step **730** combines the capabilities of the correlation network and the decompression systems. The correlation network provides the enhanced and refined latent space representations, while the decompression systems use these representations to generate the reconstructed data. In a step **740**, the restored, decompressed data set is outputted. The restored data set represents the reconstructed version of the original data, which includes recovered information lost during the compression process. The outputted data set more closely resembles the original data than would a decompressed output passed solely through a decoder network.

(116) FIG. **8** is a block diagram illustrating an exemplary system architecture for compressing and restoring multi-modal IoT data using a system for compressing and restoring data using multi-level autoencoders and correlation networks. The multi-modal IoT data organizer **800** is responsible for collecting and organizing data streams from various IoT sensors and devices. It receives raw multi-modal data from multiple sources, such as but not limited to time-series data from temperature and humidity sensors, image data from cameras, audio data from microphones, and motion data from accelerometers. The multi-modal IoT data organizer **800** may perform necessary preprocessing tasks, such as data cleaning, normalization, synchronization, and alignment of different data types, to ensure the multi-modal data is in a suitable format for further processing. The preprocessed multi-modal IoT data is then passed to a data preprocessor **810**. The data preprocessor **810** prepares the multi-modal data for compression by transforming it into a unified latent space representation. It applies techniques such as modality-specific feature extraction, dimensionality reduction, and data normalization to extract meaningful features and reduce the dimensionality of the data across all modalities. The unified latent space representation captures the essential characteristics of the multi-modal IoT data while reducing its size.

(117) The multi-layer multi-modal variational autoencoder **820** is responsible for compressing and decompressing the unified latent space representation of the multi-modal IoT data. It consists of an encoder network **821** and a decoder network **822**. The encoder network **821** takes the unified latent space representation as input and progressively compresses it through a series of modality-specific and shared layers, such as but not limited to convolutional layers for image data, recurrent layers for time-series data, and fully connected layers for numerical data. The compressed representation passes through a bottleneck layer which transforms the original multi-modal data to have a significantly reduced dimensionality in a shared latent space. Further, the encoder network **821** manages the compression process and stores the compressed representation of the multi-modal IoT data. It determines the optimal compression settings based on factors such as the desired compression ratio, data characteristics across modalities, and available storage resources. The compressed representation is efficiently stored or transmitted, reducing the storage and bandwidth requirements for multi-modal IoT data.

(118) The decoder network **822** is responsible for reconstructing the original multi-modal IoT data from the compressed representation. It utilizes the multi-layer multi-modal variational autoencoder **820** to map the compressed representation back to the original multi-modal data space. The decoder network consists of layers such as transposed convolutional layers, upsampling layers, and fully connected layers, tailored to each data modality. It learns to reconstruct the original multi-modal data by minimizing the reconstruction error between the decompressed output and the original IoT sensor data across all modalities. The decompressed output **850** represents the decompressed multi-

modal IoT data obtained from the decoder network **822**. It closely resembles the original data and retains the essential information captured by the various IoT sensors and devices, but includes some information lost during the compression process. The decompressed output **850** may be further processed, analyzed, or utilized by downstream applications or systems.

(119) To further enhance the compression and reconstruction quality, the system includes a multi-modal correlation network **830**. The multi-modal correlation network **830** learns and exploits correlations and patterns within and across the different modalities of IoT data to improve the reconstruction process. It consists of multiple correlation layers that capture dependencies and relationships among different sensors, data streams, and modalities. The multi-modal correlation network **830** helps in preserving important information that may have been lost during the compression process, particularly focusing on cross-modal relationships. Following the identification of dependencies and relationships among different data streams and modalities, the multi-modal correlation network **830** reconstructs a decompressed output **850** into a restored output **860** which recovers much of the data lost during the compression and decompression process across all modalities.

(120) The system may be trained using an end-to-end approach, where the multi-layer multi-modal variational autoencoder **820** and the multi-modal correlation network **830** are jointly optimized to minimize the reconstruction error and maximize the compression ratio across all data modalities. The training process involves feeding the multi-modal IoT data through the system, comparing the decompressed output with the original data for each modality, and updating the network parameters using backpropagation and gradient descent techniques. The proposed system offers several advantages for multi-modal IoT data compression. It achieves high compression ratios while preserving the essential information in the data across different modalities. The multi-layer multi-modal variational autoencoder **820** learns compact and meaningful representations of the data, exploiting spatial, temporal, and cross-modal correlations. The multi-modal correlation network **830** further enhances the compression quality by capturing dependencies and patterns within and across the data modalities. Moreover, the system is adaptable and can handle various types of multi-modal IoT data, making it suitable for a wide range of IoT applications. It can be deployed on resource-constrained IoT devices or edge servers, reducing storage and transmission costs while maintaining data quality across all modalities.

(121) FIG. **9** is a flow diagram illustrating an exemplary method for compressing and decompressing multi-modal IoT data using a system for compressing and restoring data using multi-level autoencoders and correlation networks. In a first step **900**, incoming multi-modal IoT data is organized based on its data type and origin. Multi-modal IoT data can be generated from various sources, such as but not limited to temperature sensors, humidity sensors, pressure sensors, accelerometers, cameras, microphones, or any other devices deployed in an IoT network. Each data type captures specific measurements or information relevant to its function. The organization step involves categorizing and grouping the incoming multi-modal IoT data based on its modality (e.g., time-series, image, audio) and the type of device it originated from. This step helps to maintain a structured and organized representation of the diverse data types, facilitating subsequent processing and analysis.

(122) In a step **910**, the latent space vectors for each multi-modal IoT data set are preprocessed. Latent space vectors are lower-dimensional representations of the original data that capture the essential features and patterns across different modalities. Preprocessing the latent space vectors involves applying various techniques to ensure data quality, consistency, and compatibility across all data types. This may include but is not limited to data cleaning, normalization, feature scaling, or dimensionality reduction, with specific techniques applied as appropriate for each data modality. The preprocessing step aims to remove any noise, outliers, or inconsistencies in the latent space vectors and prepare them for the compression process while ensuring alignment and synchronization between different data types.

(123) Step **920** involves compressing each multi-modal IoT data set using a multi-layer multi-modal variational autoencoder network. This deep learning model is designed to learn compact and meaningful representations of the multi-modal input data. It comprises an encoder network and a decoder network. The encoder network takes the preprocessed latent space vectors as input and progressively compresses them through a series of modality-specific and shared layers, such as convolutional layers for image data, recurrent layers for time-series data, and fully connected layers for numerical data. The compressed representation is obtained at the bottleneck layer of the encoder network, which has a significantly reduced dimensionality compared to the original data and represents a unified latent space for all modalities. The multi-layer multi-modal variational autoencoder network includes a compression system that specifically handles the compression of diverse IoT data types. The compression system can employ various techniques, such as quantization, entropy coding, or sparse representations, to achieve efficient compression while preserving the essential information across all data modalities. The compression system outputs a compressed multi-modal IoT data set, which is a compact representation of the original diverse data.

(124) In step **930**, the original multi-modal IoT data is decompressed using a decoder network. The decoder network is responsible for reconstructing the original data from the compressed representation in the unified latent space. It takes the compressed multi-modal IoT data sets and applies a series of decompression operations, such as transposed convolutions or upsampling layers, tailored to each data modality, to map the compressed data back to its original dimensionality and format for each data type.

(125) In step **940**, correlations between compressed multi-modal IoT data sets are identified using a multi-modal correlation network. This separate deep learning model learns to capture the relationships and dependencies among different compressed multi-modal IoT data sets, both within and across modalities. It takes the decompressed data sets as input and identifies patterns, similarities, and correlations among them, paying particular attention to cross-modal relationships. The multi-modal correlation network can utilize techniques such as convolutional layers, attention mechanisms, or graph neural networks to effectively model the interactions and dependencies between the compressed data sets of various types. The identified correlations provide valuable insights into how different multi-modal IoT data sets are related and how they influence each other across modalities. These correlations can be used to improve the compression efficiency and enhance the restoration quality of the data.

(126) In step **950**, the multi-modal correlation network creates a restored, more reconstructed version of the decompressed output. By leveraging correlations between decompressed outputs both within and across modalities, the multi-modal correlation network is able to recover a large portion of information lost during the compression and decompression process. The restored, reconstructed output is similar to the decompressed output and the original multi-modal input, but recovers information that may have been missing in the decompressed output, providing a more accurate representation of the original diverse IoT data.

(127) FIG. **10** is a block diagram illustrating an exemplary system architecture for a subsystem of the system for compressing and restoring multi-modal data using multi-level autoencoders and correlation networks, the decompressed multi-modal output organizer. In one embodiment, the decompressed multi-modal output organizer **190** may create a matrix of n by n data sets where each data set represents a decompressed set of information from various modalities. In the embodiment depicted, the decompressed multi-modal output organizer **190** outputs a 4 by 4 matrix of decompressed multi-modal data sets. The organizer **190** may organize the decompressed multi-modal data sets into groups based on how correlated each data set is to each other, both within and across modalities. For example, decompressed multi-modal data set 1, which includes 1000a, 1000b, 1000c, and 1000n, is a set of four data sets that the decompressed multi-modal output organizer **190** has determined to be highly correlated, potentially representing different aspects or

modalities of the same phenomenon (e.g., temperature data, humidity data, image data, and audio data related to a single event). The same principle applies for decompressed multi-modal data sets 2, 3, and 4.

(128) The decompressed multi-modal output organizer primes the multi-modal correlation network **160** to receive an already organized plurality of multi-modal inputs. The multi-modal correlation network may take a plurality of decompressed multi-modal data sets as its input, depending on the size of the organized matrix produced by the decompressed multi-modal output organizer **190**. For example, in the embodiment depicted in FIG. **10**, the decompressed multi-modal output organizer **190** produces a 4 by 4 matrix of multi-modal data sets. The multi-modal correlation network in turn receives a 4-element multi-modal data set as its input. If decompressed multi-modal data set 1 were to be processed by the multi-modal correlation network **160**, the network **160** may take 1000a, 1000b, 1000c, and 1000n as the inputs and process all four multi-modal data sets together, leveraging both intra-modal and cross-modal relationships. By clustering multi-modal data sets together into groups based on how correlated they are across different modalities, the decompressed multi-modal output organizer **190** allows the multi-modal correlation network **160** to produce outputs that better encompass the original pre-compressed and decompressed multi-modal data sets. More information may be recovered by the multi-modal correlation network **160** when the inputs are already highly correlated across various data types, enabling the network to exploit synergies between different modalities and reconstruct a more accurate representation of the original multi-modal data.

(129) FIG. **11** is a flow diagram illustrating an exemplary method for organizing restored, decompressed data sets after correlation network processing. In a first step **1100**, access a plurality of restored data sets. In a step **1110**, organize the plurality of restored data sets based on similarities if necessary. In a step **1120**, output a plurality of restored, potentially organizer data sets. This method essentially reassesses the organizational grouping performed by the decompressed output organizer **190**. The correlation network **160** may output a matrix where the matrix contains a plurality of restored, decompressed data sets. The final output of the system may reorganize the restored, decompressed data sets within the outputted matrix based on user preference and the correlations between each data set within the matrix.

(130) FIG. **19** is a flow diagram illustrating an exemplary method for compressing a multi-modal data input on an edge device using a system for compressing and restoring data using variational autoencoders, according to an embodiment. In a first step **1900**, multi-modal data streams are received from edge device sensors. These data streams can comprise various types of data including but not limited to images, audio, text, time-series data, and sensor measurements. The raw data arrives in different formats and modalities, requiring careful handling to maintain data integrity across all types. For example, in an industrial setting, this might include temperature readings, vibration measurements, acoustic sensors, and machine vision inputs.

(131) In a step **1910**, lightweight preprocessing parameters are initialized based on the edge device's capabilities. This step involves assessing available computational resources, memory constraints, and processing power of the edge device to optimize the preprocessing pipeline. The system configures parameters such as batch sizes, sampling rates, and buffer sizes to ensure efficient processing while working within hardware limitations.

(132) In a step **1920**, initial preprocessing is performed on the received data streams. This includes essential data cleaning, noise reduction, and format standardization optimized for edge deployment. The preprocessing step ensures that data from different modalities is properly aligned and synchronized, while maintaining computational efficiency. For example, image data might undergo basic resizing and normalization, while time-series data might be resampled to a consistent frequency.

(133) In a step **1930**, compressed representations are generated using the edge encoder. The edge encoder network, which represents the encoding portion of the multi-modal variational

autoencoder, progressively reduces the dimensionality of the preprocessed data while preserving essential features across modalities. This compression occurs in a unified latent space that maintains relationships between different data types.

(134) In a step **1940**, compressed representations are cached in local storage on the edge device. This caching mechanism allows the system to temporarily store compressed data, enabling batch processing and helping manage network bandwidth usage. The local cache also provides resilience against network interruptions and allows for efficient data handling.

(135) In a step **1950**, homomorphic operations are executed on the cached compressed data. These operations enable computations to be performed directly on the compressed representations without requiring decompression. This capability is useful for maintaining data privacy and reducing computational overhead, as it allows for meaningful data processing while keeping information in its compressed form.

(136) In a step **1960**, the processed compressed data is synchronized with the cloud service. This synchronization step involves transmitting the compressed and processed data to cloud infrastructure for further processing, decompression, and analysis. The system manages this transfer efficiently, potentially using batch transmission to optimize network usage and ensure reliable data delivery.

(137) FIG. **20** is a flow diagram illustrating an exemplary method for decompressing compressed multi-modal data input on using a system for compressing and restoring data using variational autoencoders stores in a cloud environment, according to an embodiment. In a first step **2000**, synchronized compressed data is received from edge devices in the cloud infrastructure. This represents the compressed representations that have been processed and transmitted by various edge devices, potentially including multiple data modalities from different sources. The synchronization ensures that related data from different edge devices is properly aligned in time and context.

(138) In a step **2010**, the integrity of received compressed representations is validated. This validation process ensures that the compressed data maintains its mathematical properties and hasn't been corrupted during transmission. The system verifies that the compressed representations remain in the correct format for the unified latent space and checks for any anomalies or inconsistencies in the data structure.

(139) In a step **2020**, the compressed data is processed in the unified latent space. The unified latent space provides a continuous and differentiable representation where all data modalities can be processed together while maintaining their relationships. This step may involve aligning different compressed representations and preparing them for global operations.

(140) In a step **2030**, global homomorphic operations are executed across the aggregated data. These operations allow for computations to be performed on the compressed data from multiple edge devices simultaneously, without requiring decompression. This capability is particularly valuable for analyzing patterns across multiple data sources while maintaining data privacy and computational efficiency.

(141) In a step **2040**, the processed data is decompressed using the cloud decoder. The cloud decoder network maps the processed compressed representations back to their original data spaces while preserving the relationships between different modalities. This step begins the reconstruction process of the original multi-modal data.

(142) In a step **2050**, decompressed outputs are organized using a correlation network. The correlation network analyzes relationships both within and across modalities to enhance the reconstruction quality. It leverages these identified patterns to recover information that may have been lost during compression and helps establish connections between different data types.

(143) In a step **2060**, the final reconstructed multi-modal output is generated. This output represents the fully processed and reconstructed version of the original data, enhanced through the correlation network's pattern recognition and relationship modeling. The final output maintains the essential

features and relationships present in the original multi-modal data while benefiting from the global processing performed in the cloud infrastructure.

(144) Exemplary Computing Environment

(145) FIG. 12 illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part. This exemplary computing environment describes computer-related components and processes supporting enabling disclosure of computer-implemented embodiments. Inclusion in this exemplary computing environment of well-known processes and computer components, if any, is not a suggestion or admission that any embodiment is no more than an aggregation of such processes or components. Rather, implementation of an embodiment using processes and components described in this exemplary computing environment will involve programming or configuration of such processes and components resulting in a machine specially programmed or configured for such implementation. The exemplary computing environment described herein is only one example of such an environment and other configurations of the components and processes are possible, including other relationships between and among components, and/or absence of some processes or components described. Further, the exemplary computing environment described herein is not intended to suggest any limitation as to the scope of use or functionality of any embodiment implemented, in whole or in part, on components or processes described herein.

(146) The exemplary computing environment described herein comprises a computing device **10** (further comprising a system bus **11**, one or more processors **20**, a system memory **30**, one or more interfaces **40**, one or more non-volatile data storage devices **50**), external peripherals and accessories **60**, external communication devices **70**, remote computing devices **80**, and cloud-based services **90**. The system memory **30** can include a computer readable storage medium.

(147) System bus **11** couples the various system components, coordinating operation of and data transmission between those various system components. System bus **11** represents one or more of any type or combination of types of wired or wireless bus structures including, but not limited to, memory busses or memory controllers, point-to-point connections, switching fabrics, peripheral busses, accelerated graphics ports, and local busses using any of a variety of bus architectures. By way of example, such architectures include, but are not limited to, Industry Standard Architecture (ISA) busses, Micro Channel Architecture (MCA) busses, Enhanced ISA (EISA) busses, Video Electronics Standards Association (VESA) local busses, a Peripheral Component Interconnects (PCI) busses also known as a Mezzanine busses, or any selection of, or combination of, such busses. Depending on the specific physical implementation, one or more of the processors **20**, system memory **30** and other components of the computing device **10** can be physically co-located or integrated into a single physical component, such as on a single chip. In such a case, some or all of system bus **11** can be electrical pathways within a single chip structure.

(148) Computing device may further comprise externally-accessible data input and storage devices **12** such as compact disc read-only memory (CD-ROM) drives, digital versatile discs (DVD), or other optical disc storage for reading and/or writing optical discs **62**; magnetic cassettes, magnetic tape, magnetic disk storage, or other magnetic storage devices; or any other medium which can be used to store the desired content and which can be accessed by the computing device **10**.

Computing device may further comprise externally-accessible data ports or connections **12** such as serial ports, parallel ports, universal serial bus (USB) ports, and infrared ports and/or transmitter/receivers. Computing device may further comprise hardware for wireless communication with external devices such as IEEE 1394 ("Firewire") interfaces, IEEE 802.11 wireless interfaces, BLUETOOTH® wireless interfaces, and so forth. Such ports and interfaces may be used to connect any number of external peripherals and accessories **60** such as visual displays, monitors, and touch-sensitive screens **61**, USB solid state memory data storage drives (commonly known as "flash drives" or "thumb drives") **63**, printers **64**, pointers and manipulators such as mice **65**, keyboards **66**, and other devices **67** such as joysticks and gaming pads, touchpads,

additional displays and monitors, and external hard drives (whether solid state or disc-based), microphones, speakers, cameras, and optical scanners.

(149) Processors **20** are logic circuitry capable of receiving programming instructions and processing (or executing) those instructions to perform computer operations such as retrieving data, storing data, and performing mathematical calculations. Processors **20** are not limited by the materials from which they are formed or the processing mechanisms employed therein, but are typically comprised of semiconductor materials into which many transistors are formed together into logic gates on a chip (i.e., an integrated circuit or IC). The term processor includes any device capable of receiving and processing instructions including, but not limited to, processors operating on the basis of quantum computing, optical computing, mechanical computing (e.g., using nanotechnology entities to transfer data), and so forth. Depending on configuration, computing device **10** may comprise more than one processor. For example, computing device **10** may comprise one or more central processing units (CPUs) **21**, each of which itself has multiple processors or multiple processing cores, each capable of independently or semi-independently processing programming instructions. Further, computing device **10** may comprise one or more specialized processors such as a graphics processing unit (GPU) **22** configured to accelerate processing of computer graphics and images via a large array of specialized processing cores arranged in parallel.

(150) System memory **30** is processor-accessible data storage in the form of volatile and/or nonvolatile memory. System memory **30** may be either or both of two types: non-volatile memory and volatile memory. Non-volatile memory **30a** is not erased when power to the memory is removed, and includes memory types such as read only memory (ROM), electronically-erasable programmable memory (EEPROM), and rewritable solid-state memory (commonly known as “flash memory”). Non-volatile memory **30a** is typically used for long-term storage of a basic input/output system (BIOS) **31**, containing the basic instructions, typically loaded during computer startup, for transfer of information between components within computing device, or a unified extensible firmware interface (UEFI), which is a modern replacement for BIOS that supports larger hard drives, faster boot times, more security features, and provides native support for graphics and mouse cursors. Non-volatile memory **30a** may also be used to store firmware comprising a complete operating system **35** and applications **36** for operating computer-controlled devices. The firmware approach is often used for purpose-specific computer-controlled devices such as appliances and Internet-of-Things (IoT) devices where processing power and data storage space is limited. Volatile memory **30b** is erased when power to the memory is removed and is typically used for short-term storage of data for processing. Volatile memory **30b** includes memory types such as random-access memory (RAM), and is normally the primary operating memory into which the operating system **35**, applications **36**, program modules **37**, and application data **38** are loaded for execution by processors **20**. Volatile memory **30b** is generally faster than non-volatile memory **30a** due to its electrical characteristics and is directly accessible to processors **20** for processing of instructions and data storage and retrieval. Volatile memory **30b** may comprise one or more smaller cache memories which operate at a higher clock speed and are typically placed on the same IC as the processors to improve performance.

(151) Interfaces **40** may include, but are not limited to, storage media interfaces **41**, network interfaces **42**, display interfaces **43**, and input/output interfaces **44**. Storage media interface **41** provides the necessary hardware interface for loading data from non-volatile data storage devices **50** into system memory **30** and storage data from system memory **30** to non-volatile data storage device **50**. Network interface **42** provides the necessary hardware interface for computing device **10** to communicate with remote computing devices **80** and cloud-based services **90** via one or more external communication devices **70**. Display interface **43** allows for connection of displays **61**, monitors, touchscreens, and other visual input/output devices. Display interface **43** may include a graphics card for processing graphics-intensive calculations and for handling demanding display

requirements. Typically, a graphics card includes a graphics processing unit (GPU) and video RAM (VRAM) to accelerate display of graphics. One or more input/output (I/O) interfaces **44** provide the necessary support for communications between computing device **10** and any external peripherals and accessories **60**. For wireless communications, the necessary radio-frequency hardware and firmware may be connected to I/O interface **44** or may be integrated into I/O interface **44**.

(152) Non-volatile data storage devices **50** are typically used for long-term storage of data. Data on non-volatile data storage devices **50** is not erased when power to the non-volatile data storage devices **50** is removed. Non-volatile data storage devices **50** may be implemented using any technology for non-volatile storage of content including, but not limited to, CD-ROM drives, digital versatile discs (DVD), or other optical disc storage; magnetic cassettes, magnetic tape, magnetic disc storage, or other magnetic storage devices; solid state memory technologies such as EEPROM or flash memory; or other memory technology or any other medium which can be used to store data without requiring power to retain the data after it is written. Non-volatile data storage devices **50** may be non-removable from computing device **10** as in the case of internal hard drives, removable from computing device **10** as in the case of external USB hard drives, or a combination thereof, but computing device will typically comprise one or more internal, non-removable hard drives using either magnetic disc or solid-state memory technology. Non-volatile data storage devices **50** may store any type of data including, but not limited to, an operating system **51** for providing low-level and mid-level functionality of computing device **10**, applications **52** for providing high-level functionality of computing device **10**, program modules **53** such as containerized programs or applications, or other modular content or modular programming, application data **54**, and databases **55** such as relational databases, non-relational databases, object oriented databases, BOSQL databases, and graph databases.

(153) Applications (also known as computer software or software applications) are sets of programming instructions designed to perform specific tasks or provide specific functionality on a computer or other computing devices. Applications are typically written in high-level programming languages such as C++, Java, and Python, which are then either interpreted at runtime or compiled into low-level, binary, processor-executable instructions operable on processors **20**. Applications may be containerized so that they can be run on any computer hardware running any known operating system. Containerization of computer software is a method of packaging and deploying applications along with their operating system dependencies into self-contained, isolated units known as containers. Containers provide a lightweight and consistent runtime environment that allows applications to run reliably across different computing environments, such as development, testing, and production systems.

(154) The memories and non-volatile data storage devices described herein do not include communication media. Communication media are means of transmission of information such as modulated electromagnetic waves or modulated data signals configured to transmit, not store, information. By way of example, and not limitation, communication media includes wired communications such as sound signals transmitted to a speaker via a speaker wire, and wireless communications such as acoustic waves, radio frequency (RF) transmissions, infrared emissions, and other wireless media.

(155) External communication devices **70** are devices that facilitate communications between computing device and either remote computing devices **80**, or cloud-based services **90**, or both. External communication devices **70** include, but are not limited to, data modems **71** which facilitate data transmission between computing device and the Internet **75** via a common carrier such as a telephone company or internet service provider (ISP), routers **72** which facilitate data transmission between computing device and other devices, and switches **73** which provide direct data communications between devices on a network. Here, modem **71** is shown connecting computing device **10** to both remote computing devices **80** and cloud-based services **90** via the Internet **75**. While modem **71**, router **72**, and switch **73** are shown here as being connected to network interface

42, many different network configurations using external communication devices **70** are possible. Using external communication devices **70**, networks may be configured as local area networks (LANs) for a single location, building, or campus, wide area networks (WANs) comprising data networks that extend over a larger geographical area, and virtual private networks (VPNs) which can be of any size but connect computers via encrypted communications over public networks such as the Internet **75**. As just one exemplary network configuration, network interface **42** may be connected to switch **73** which is connected to router **72** which is connected to modem **71** which provides access for computing device **10** to the Internet **75**. Further, any combination of wired **77** or wireless **76** communications between and among computing device **10**, external communication devices **70**, remote computing devices **80**, and cloud-based services **90** may be used. Remote computing devices **80**, for example, may communicate with computing device through a variety of communication channels **74** such as through switch **73** via a wired **77** connection, through router **72** via a wireless connection **76**, or through modem **71** via the Internet **75**. Furthermore, while not shown here, other hardware that is specifically designed for servers may be employed. For example, secure socket layer (SSL) acceleration cards can be used to offload SSL encryption computations, and transmission control protocol/internet protocol (TCP/IP) offload hardware and/or packet classifiers on network interfaces **42** may be installed and used at server devices.

(156) In a networked environment, certain components of computing device **10** may be fully or partially implemented on remote computing devices **80** or cloud-based services **90**. Data stored in non-volatile data storage device **50** may be received from, shared with, duplicated on, or offloaded to a non-volatile data storage device on one or more remote computing devices **80** or in a cloud computing service **92**. Processing by processors **20** may be received from, shared with, duplicated on, or offloaded to processors of one or more remote computing devices **80** or in a distributed computing service **93**. By way of example, data may reside on a cloud computing service **92**, but may be usable or otherwise accessible for use by computing device **10**. Also, certain processing subtasks may be sent to a microservice **91** for processing with the result being transmitted to computing device **10** for incorporation into a larger processing task. Also, while components and processes of the exemplary computing environment are illustrated herein as discrete units (e.g., OS **51** being stored on non-volatile data storage device **51** and loaded into system memory **35** for use) such processes and components may reside or be processed at various times in different components of computing device **10**, remote computing devices **80**, and/or cloud-based services **90**.

(157) In an implementation, the disclosed systems and methods may utilize, at least in part, containerization techniques to execute one or more processes and/or steps disclosed herein. Containerization is a lightweight and efficient virtualization technique that allows you to package and run applications and their dependencies in isolated environments called containers. One of the most popular containerization platforms is Docker, which is widely used in software development and deployment. Containerization, particularly with open-source technologies like Docker and container orchestration systems like Kubernetes, is a common approach for deploying and managing applications. Containers are created from images, which are lightweight, standalone, and executable packages that include application code, libraries, dependencies, and runtime. Images are often built from a Dockerfile or similar, which contains instructions for assembling the image. Dockerfiles are configuration files that specify how to build a Docker image. Systems like Kubernetes also support containers or CRI-O. They include commands for installing dependencies, copying files, setting environment variables, and defining runtime configurations. Docker images are stored in repositories, which can be public or private. Docker Hub is an exemplary public registry, and organizations often set up private registries for security and version control using tools such as Hub, JFrog Artifactory and Bintray, Github Packages or Container registries. Containers can communicate with each other and the external world through networking. Docker provides a bridge network by default, but can be used with custom networks. Containers within the same

network can communicate using container names or IP addresses.

(158) Remote computing devices **80** are any computing devices not part of computing device **10**. Remote computing devices **80** include, but are not limited to, personal computers, server computers, thin clients, thick clients, personal digital assistants (PDAs), mobile telephones, watches, tablet computers, laptop computers, multiprocessor systems, microprocessor based systems, set-top boxes, programmable consumer electronics, video game machines, game consoles, portable or handheld gaming units, network terminals, desktop personal computers (PCs), minicomputers, main frame computers, network nodes, virtual reality or augmented reality devices and wearables, and distributed or multi-processing computing environments. While remote computing devices **80** are shown for clarity as being separate from cloud-based services **90**, cloud-based services **90** are implemented on collections of networked remote computing devices **80**.

(159) Cloud-based services **90** are Internet-accessible services implemented on collections of networked remote computing devices **80**. Cloud-based services are typically accessed via application programming interfaces (APIs) which are software interfaces which provide access to computing services within the cloud-based service via API calls, which are pre-defined protocols for requesting a computing service and receiving the results of that computing service. While cloud-based services may comprise any type of computer processing or storage, three common categories of cloud-based services **90** are microservices **91**, cloud computing services **92**, and distributed computing services **93**.

(160) Microservices **91** are collections of small, loosely coupled, and independently deployable computing services. Each microservice represents a specific computing functionality and runs as a separate process or container. Microservices promote the decomposition of complex applications into smaller, manageable services that can be developed, deployed, and scaled independently. These services communicate with each other through well-defined application programming interfaces (APIs), typically using lightweight protocols like HTTP, gRPC, or message queues such as Kafka. Microservices **91** can be combined to perform more complex processing tasks.

(161) Cloud computing services **92** are delivery of computing resources and services over the Internet **75** from a remote location. Cloud computing services **92** provide additional computer hardware and storage on as-needed or subscription basis. Cloud computing services **92** can provide large amounts of scalable data storage, access to sophisticated software and powerful server-based processing, or entire computing infrastructures and platforms. For example, cloud computing services can provide virtualized computing resources such as virtual machines, storage, and networks, platforms for developing, running, and managing applications without the complexity of infrastructure management, and complete software applications over the Internet on a subscription basis.

(162) Distributed computing services **93** provide large-scale processing using multiple interconnected computers or nodes to solve computational problems or perform tasks collectively. In distributed computing, the processing and storage capabilities of multiple machines are leveraged to work together as a unified system. Distributed computing services are designed to address problems that cannot be efficiently solved by a single computer or that require large-scale computational power. These services enable parallel processing, fault tolerance, and scalability by distributing tasks across multiple nodes.

(163) As can now be appreciated, disclosed embodiments utilize a variational autoencoder to enable homomorphic compression techniques. Input data is compressed into a latent space using an encoder network of a variational autoencoder. Homomorphic operations are performed on the compressed data in the latent space. The latent space compressed data is decompressed using a decoder network of the variational autoencoder. The homomorphic operations can enable performing operations while the data is in a compressed form, and preserving the results of those operations after the data is in a decompressed form. One or more embodiments can compress data using partially homomorphic encryption (PHE). The PHE of disclosed embodiments allows for

specific types of computations to be performed on encrypted data without decrypting it. For example, in a partially homomorphic encryption scheme that supports addition and multiplication, embodiments can include encrypting a value x , encrypting a threshold value T , and then performing a homomorphic comparison operation to check if $x < T$. In a medical records application, the value T can be a year, and the value x can be a birth year of a patient in a medical record. The result of this operation provides an encrypted value that represents the comparison result. Continuing with the example, disclosed embodiments can enable checking if a patient referred to in a medical record is above a certain age, without needed to decrypt/decompress the medical record. This powerful feature enables third-party verification operations. Continuing with the example, a third party can perform age verification on medical records while in an encrypted form, such that the third party cannot view any personally sensitive identifying information in the medical records. Thus, the homomorphic encryption and compression of disclosed embodiments can provide a range of benefits including enhanced privacy, data security, reduced storage requirements, efficient data processing, and secure data outsourcing.

(164) Although described above as a physical device, referring again to FIG. 12, computing device 10 can be a virtual computing device, in which case the functionality of the physical components herein described, such as processors 20, system memory 30, network interfaces 40, and other like components can be provided by computer-executable instructions. Such computer-executable instructions can execute on a single physical computing device, or can be distributed across multiple physical computing devices, including being distributed across multiple physical computing devices in a dynamic manner such that the specific, physical computing devices hosting such computer-executable instructions can dynamically change over time depending upon need and availability. In the situation where computing device 10 is a virtualized device, the underlying physical computing devices hosting such a virtualized computing device can, themselves, comprise physical components analogous to those described above, and operating in a like manner. Furthermore, virtual computing devices can be utilized in multiple layers with one virtual computing device executing within the construct of another virtual computing device. Thus, computing device 10 may be either a physical computing device or a virtualized computing device within which computer-executable instructions can be executed in a manner consistent with their execution by a physical computing device. Similarly, terms referring to physical components of the computing device, as utilized herein, mean either those physical components or virtualizations thereof performing the same or equivalent functions.

(165) The skilled person will be aware of a range of possible modifications of the various aspects described above. Accordingly, the present invention is defined by the claims and their equivalents.

Claims

1. A computer system comprising a hardware memory, wherein the computer system is configured to execute software instructions stored on nontransitory machine-readable storage media that: preprocess raw multi-modal data at an edge computing device to generate a plurality of input data sets; compress the plurality of input data sets into a plurality of compressed data sets using an encoder within a multi-modal variational autoencoder at the edge computing device; transmit the compressed data sets from the edge computing device to a cloud computing device; and decompress the plurality of compressed data sets using a decoder located within the cloud computing device to obtain a plurality of reconstructed data sets.
2. The computer system of claim 1, wherein the plurality of reconstructed data sets are organized by a correlation network, wherein the correlation network recovers a plurality of information lost during compression by the edge device.
3. The computer system of claim 1, wherein the multi-modal data comprises at least two of: image data, audio data, text data, time-series data, and sensor data.

4. The computer system of claim 1, wherein the computer system is further configured to execute software instructions stored on nontransitory machine-readable storage media that: use the multi-modal variational autoencoder to perform partially homomorphic compression in a unified latent space representing all data modalities.

5. A method for compressing and restoring multi-modal data, comprising steps of: preprocessing raw multi-modal data at an edge computing device to generate a plurality of input data sets; compressing the plurality of input data sets into a plurality of compressed data sets using an encoder within a multi-modal variational autoencoder at the edge computing device; transmitting the compressed data sets from the edge computing device to a cloud computing device; and decompressing the plurality of compressed data sets using a decoder located within the cloud computing device to obtain a plurality of reconstructed data sets.

6. The method of claim 5, wherein the plurality of reconstructed data sets are organized by a correlation network, wherein the correlation network recovers a plurality of information lost during compression by the edge device.

7. The method of claim 5, wherein the multi-modal data comprises at least two of: image data, audio data, text data, time-series data, and sensor data.

8. The method of claim 5, further comprising using the multi-modal variational autoencoder to perform partially homomorphic compression in a unified latent space representing all data modalities.
