

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417177
Kind Code	B1
Date of Patent	September 16, 2025
Inventor(s)	Lei; Ming Liang Milton et al.

Systems and methods for multicasting data

Abstract

A first instruction to load a first data into a first register file associated with a first subpartition unit is received. A second instruction to load the first data into a second register file associated with a second subpartition unit is received. The first instruction and the second instruction are coalesced into a first entry of a request coalescer based on instruction identifiers. The first entry is associated with the first data. Responsive to a determination that the first data is available in the cache, the first data is multicast from the cache to the first register file and the second register file.

Inventors: Lei; Ming Liang Milton (Santa Clara, CA), Schottmiller; Jeffery Michael (Raleigh, NC), Patel; Manan (San Jose, CA), Ghoshal; Pritha (Apex, NC), Fetterman; Michael Alan (Lancaster, MA), Ohannessian, Jr.; Robert (Austin, TX), Kaushik; Praveen Kumar (Karnataka, IN)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 1000007781783

Assignee: NVIDIA Corporation (Santa Clara, CA)

Appl. No.: 18/604184

Filed: March 13, 2024

Publication Classification

Int. Cl.: G06F12/0804 (20160101); G06F12/0868 (20160101); G06F12/0875 (20160101)

U.S. Cl.:

CPC G06F12/0804 (20130101); G06F12/0868 (20130101); G06F12/0875 (20130101);

Field of Classification Search

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11507506	12/2021	Kipp	N/A	G06F 12/0804
2018/0302468	12/2017	Hu	N/A	H04L 67/565

Primary Examiner: Ruiz; Aracelis

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

FIELD OF THE INVENTION

(1) Embodiments of the present disclosure generally relate to parallel processing systems. Specifically, embodiments of the present disclosure relate to parallel processing systems and methods for multicasting data.

BACKGROUND

(2) Parallel processing in high-performance computing (HPC) systems involve the simultaneous execution of multiple computational tasks or operations. This is done by breaking down larger computations into smaller, independent, subtasks that be processed concurrently by multiprocessors. In some instances, parallel processing involves distributed computing, such that task are distributed across multiple computing clusters. Each cluster may operate independently, and communication can be facilitated to share results.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 illustrates a parallel processing unit, in accordance with at least one embodiment of the present disclosure.
- (2) FIG. 2 illustrates a conceptual diagram as instructions requesting constant data are coalesced and constant data is multicast to associated uniform register files, in accordance with at least one embodiment of the present disclosure.
- (3) FIG. 3A illustrates an example request coalescer Content Addressable Memory (CAM) table, in accordance with at least one embodiment of the present disclosure.
- (4) FIG. 3B illustrates example pseudocode to coalesce a request into the example request coalescer CAM table of FIG. 3A, in accordance with at least one embodiment of the present disclosure.
- (5) FIG. 3C illustrates a continuation of the example pseudocode to coalesce the request into the example request coalescer CAM table of FIG. 3A, in accordance with at least one embodiment of the present disclosure.
- (6) FIG. 4 illustrates a flowchart of a method for multicasting data, in accordance with at least one embodiment of the present disclosure.
- (7) FIG. 5 illustrates a parallel processing unit, in accordance with an embodiment.

- (8) FIG. 6A illustrates a general processing cluster within the parallel processing unit of FIG. 3, in accordance with an embodiment.
- (9) FIG. 6B illustrates a memory partition unit of the parallel processing unit of FIG. 5, in accordance with an embodiment.
- (10) FIG. 7A illustrates the streaming multi-processor of FIG. 6A, in accordance with an embodiment.
- (11) FIG. 7B is a conceptual diagram of a processing system implemented using the PPU of FIG. 5, in accordance with an embodiment.
- (12) FIG. 7C illustrates an exemplary system in which the various architecture and/or functionality of the various previous embodiments may be implemented.
- (13) FIG. 8 is a conceptual diagram of a graphics processing pipeline implemented by the PPU of FIG. 5, in accordance with an embodiment.

DETAILED DESCRIPTION

(14) HPC systems may use specialized hardware architectures, such as parallel processing units (PPUs), to enhance parallel processing abilities. PPUs are designed to extract high performance using a large number of small, parallel execution threads on dedicated programmable multiprocessors. In PPUs, a group of threads, such as a warp, may execute the same instruction (e.g., the same instruction executed by different threads) concurrently on a multiprocessor with different input data. This execution model is referred to as Single Instruction, Multiple Thread (SIMT) and is commonly utilized in parallel computing. PPUs are designed to execute a program (e.g., a kernel, a shader program, etc.) in parallel by many warps on the PPU in which each thread of a warp typically operates on a different portion of data. In some instances, more than one multiprocessor may request the same data at different points of executing the program. To account for these situations, some PPUs may include one or more specialized memory spaces, such as a constant cache and associated uniform registers, that are optimized for storing and accessing constant data. The constant cache may be a read-only memory space configured to store constant data that remains the same across multiple warps during execution of a program. Constant data can include mathematical constants, program parameters, transformation matrices, coefficients, and any other data that remains constant during program execution. Constant data may be read from the constant cache and written into uniform registers (referred to as “writeback” herein) associated one or more threads across warps requesting the constant data. However, servicing multiple threads requesting the same constant data can create a bandwidth issue. For example, the constant cache may have a writeback bandwidth of 8 bytes per clock cycle. A first, second, third, and fourth thread may request the same 32 bytes of constant data. Each of the requests may be serviced individually for a total writeback bandwidth of 16 clock cycles.

(15) To mitigate the bandwidth associated with writeback of constant data, some conventional techniques may match addresses of incoming requests for constant data. Matching addresses can be coalesced, and associated constant data can be multicast to corresponding uniform registers. Coalescing is the process of combining multiple memory access requests, such as request to read data from a constant cache, into a single memory access requests. In the context of the present disclosure, multicasting refers to the process of sending constant data from a constant cache to multiple requestors simultaneously for writeback into uniform registers. Uniform registers can generally refer to on-chip memory structures to store variables or values that remain the same across a group of parallel threads. To match addresses among incoming requests, some conventional systems may implement wide address matching logic with an associated area cost. For example, the system may implement address matching logic to match 64-bit memory addresses. Area allocated for address matching logic can be significant for a constant cache that is heavily replicated by the count of multiprocessors on the PPU. Area is an important consideration in PPU design as it is a factor in cost efficiency, power efficiency, heat dissipation, and the like.

(16) In another conventional solution, a single thread may programmatically orchestrate

multicasting of data. The PPU may implement software to coordinate among threads prior to memory read instruction at issue. Designing such software can be challenging and introduce significant overhead in terms of design complexity.

(17) Another conventional solution may implement mandatory aging of requests to writeback constant data from the constant cache into uniform registers in order to widen a window in which writeback requests may be coalesced. However, using mandatory aging to coalesce writeback requests can increase a static latency for thread execution. Additionally, after coalescing, the data still must be broadcast to uniform registers of associated multiprocessor cores to achieve the desired writeback bandwidth reduction. Conventional techniques may blindly broadcast data for writeback to the requestors at the same time to achieve the desired reduced writeback bandwidth. Such a solution requires all requestors to have free cycle for uniform register writeback at the same time, and coordinating time of writeback with each uniform register associated with various multiprocessors can add complexity.

(18) Aspects and implementations of the present disclosure address the above deficiencies and other deficiencies of conventional PPUs by providing a technique that enables a PPU to coalesce writeback requests using a reduced number of bits. Constant data can be multicast to corresponding uniform register files of disjoint cores of a multiprocessor. A PPU can integrate several independent multiprocessors, such as streaming multiprocessors (SMs) in a single die. Each multiprocessor may include a request coalescer and a constant cache accessible to all subpartition units and associated processing cores disposed on the multiprocessor. Each subpartition unit can execute a warp, allowing for parallel execution of warps. Subpartition units can issue writeback requests to the request coalescer to read data from the constant cache and write the data to one or more uniform registers associated with the subpartition. In some embodiments, the writeback request is associated with an instruction to perform an operation using the requested data. Accordingly, the terms request and instruction may be used interchangeably herein by way of convenience, noting that aspects and implementations of the present disclosure may generally apply to coalescing requests to access constant data. In some embodiments, the writeback requests can be coalesced by the request coalescer according to an instruction identifier. When the data is available in the cache, the request coalescer can cause the data to be multicast to uniform registers associated with requests having the same instruction identifier. In some embodiments, the instruction identifier can include a cacheline identifier and a byte offset. Accordingly, the instruction identifier used to match writeback requests can be reduced from a memory address (e.g., 64 bits) to a cacheline identifier (e.g., 9 bits) and a byte offset (e.g., 16 bits).

(19) In at least one embodiment, to multicast the constant data, the request coalescer may simultaneously write the data to one or more writeback buffers. The writeback buffers can store the data until corresponding uniform register files are available for writeback. For example, a first request and a second request may request the same constant data. The request coalescer may coalesce the first request and the second request into a same entry of the request coalescer. Responsive to the constant data being available in the constant cache, the request coalescer can write the constant data to a first writeback buffer associated with the first request and a second writeback buffer associated with the second request. When a first uniform register corresponding to the first request is available for writeback, the first writeback buffer can load the constant data into the first uniform register. When a second uniform register corresponding to the second request is available, the second writeback buffer can load the constant data into the second uniform register. As such, temporarily storing constant data in generic writeback buffers can allow multicasting of constant data to separate cores of a multiprocessor without synchronizing uniform registers for writeback.

(20) Advantages of the technology disclosure herein include, but are not limited, decreased bandwidth associated with writeback of constant data into uniform register files. This can be accomplished by coalescing multiple writeback requests using minimal bits for request/instruction

identification and matching. Additionally, leveraging a generic writeback buffer for multicasting data can negate complexity and hardware associated with synchronized uniform register writeback. Accordingly, aspects and implementations of the present disclosure can provide flexibility with comparatively little complexity when coalescing writeback requests and multicasting data to uniform registers to reduce writeback bandwidth.

(21) It should be noted that various aspects of the above referenced methods and systems are described in detail herein below by way of example, rather than by way of limitation. The embodiments and examples provided below may reference writeback requests/instructions associated with particular subpartition units of a multiprocessor. However, writeback requests can generally be associated with a thread executing on any processing core(s) associated with a PPU. Additionally, the embodiments and example provided below may reference uniform register associated with particular subpartition units. However, uniform registers can be associated with a multiprocessor, a subpartition unit, an individual thread, or any processing unit of a PPU. Further, it is appreciated that a warp is used herein by way of example, and not by way of limitation, noting that aspects and implementations of the present disclosure can generally be applied to groups of threads that are scheduled together. Additionally, it is appreciated that uniform registers are used herein by way of example, and not by way of limitation. Noting that aspects and implementations of the present disclosure can utilize other registers for multicasting data. For example, data can be multicast to a general-purpose register file according to techniques described herein.

(22) FIG. 1 illustrates a parallel processing unit (PPU) **100**, in accordance with at least one embodiment of the present disclosure. The PPU **100**, for example, may include a graphical processing unit (GPU). As illustrated in FIG. 1, the PPU **100** includes one or more multiprocessor(s) **102**. Each multiprocessor **102**, for example, may be a streaming multiprocessor (SM), a compute unit (CU), a many integrated core (MIC), and the like. Each of the multiprocessors **102** includes, without limitation, multiple subpartition units **104A** through **104N** (referred to generally as “subpartition unit(s) **104**” herein), a request coalescer **120**, and a constant cache **110**.

(23) Warps assigned to the multiprocessor **102** are distributed between the subpartition units **104**. Each of the subpartition units **104** may be assigned any number of warps, however, a given warp may only be assigned to one subpartition unit **104**. As illustrated, each subpartition unit **104** may include, without limitation, multiple functional units **106A** through **106N** (referred to generally as “functional unit(s) **106**” herein), a uniform register file (URF) **108A** through **108N** (referred to generally as “URF(s) **108**” herein), and a writeback buffer **109A** through **109N** (referred to generally as “WB Buffer(s) **109**” herein). Each of the functional units **106** may include multiple processing cores (e.g., Compute Unified Device Architecture (CUDA) cores) to execute separate threads, allowing for parallel processing. In at least one embodiment, the functional units **106** may operate according to a Single-Instruction, Multiple Thread (SIMT) architecture. For example, all threads within a warp execute the same instruction at the same time, but may operate on different data.

(24) In at least one embodiment, the functional units **106** may perform a variation of computations, such as floating point 32-bit precision arithmetic, floating point 16-bit precision arithmetic, integer arithmetic of varying precisions, and the like. In at least one embodiment, the functional units **106** may include one or more tensor cores (not illustrated) to perform matrix multiplication operations using mixed-precision arithmetic. For example, the one or more tensor cores may operate on one low-precision formats, such as half-precision floating-point (FP16) to accelerate computation.

(25) As indicated above, each of the subpartition units **104** may include a uniform register file **108**. The uniform register file **108** is collections of registers used to store temporary data, intermediate results, operands during computation, and/or constant data. The uniform register file (URF) **108** is a shared memory resource that includes, without limitation, any number of uniform registers and any number of read and/or write ports. Each uniform register in the URF **108** is independently

accessible to all threads included in a warp executing on a corresponding subpartition unit **104**. The URF **108** may be implemented in any technically feasible manner. In an alternative embodiment, the uniform registers included in the URF **108** may be arranged and assigned to threads and/or warps in any technically feasible manner. In at least one embodiment, each thread running on a given functional unit **104** may use a dedicated set of registers separate from the URF **108** associated with the thread. In at least one embodiment, these registers may be private to the thread, providing a dedicated space for storing data associated with the thread.

(26) It is appreciated that the number of multiprocessors **102** in the PPU **100**, the number of subpartition units **104** in a single multiprocessor **102**, and the number of functional units **106** in a subpartition unit **104** are not limited. Each subpartition unit **104** in a multiprocessor **102** may include a URF **108** and may have access to the constant cache **110** for that multiprocessor **102**. In at least one embodiment, as in the embodiment illustrated in FIG. **1**, the constant cache **110** may be a part of a shared/local memory associated with that multiprocessor **102**. In some other embodiments, the constant cache **110** may be separate from a shared memory associated with that multiprocessor **102**.

(27) In some embodiments, the constant cache **110** may be used to store constants or values that do not change during execution of a kernel or shader program. These constants, for example, may include parameters, coefficients, or any other data that remains constant across a multiple warps. Accordingly, the constant cache **110** may be a read-only cache located closer to processing units/cores than global memory, providing faster access times for threads that access the same constant data. In some embodiments, values stored within the constant cache **110** may be updated. For example, a kernel/shader program may update a value of a global memory address associated with the constant cache and issue invalidates to the constant cache **110** to allow the constant cache **110** to observe updated values. In some embodiments, the constant cache **110** may be located on a main structural unit (e.g., a Graphics Processing Cluster (GPC)) of the PPU **100** and associated with multiple multiprocessors **102** (e.g., SMs). In some embodiments, the constant cache **110** may be associated with a single multiprocessor **102**, as illustrated with respect to FIG. **1**. For example, each constant cache **110** may be associated with a single multiprocessor **102**. The constant cache **110** may generally refer to a caching mechanism associated with a constant memory across an entire memory hierarchy of the PPU **100**.

(28) In at least one embodiment, the multiprocessors **102** may access a global memory (not illustrated) that is external to the PPU **100** through a global memory interface. The global memory may include a hierarchical cache memory (e.g., L2 cache and/or L3 cache) and a dynamic random-access memory (DRAM). In at least one embodiment, the global memory may include a memory management unit (MMU), an X-Bar or hierarchical cross-bar interconnect network, a memory partition unit, and/or memory described with reference to FIGS. **6**, **7A**, and **7B**.

(29) Multiple functional units **104** in each of the multiprocessors **102** are configured to process a plurality of threads in parallel. A thread (e.g., a thread of execution) is an instantiation of a set of instructions or a kernel configured to be executed by the functional units **104** on a particular data set. Threads of a thread block may be executed concurrently, and multiple thread blocks may be executed concurrently. In some embodiments, single-instruction multiple data (SIMD) instruction issue techniques are used to support parallel execution of a large number of threads without providing multiple independent instruction units. In other embodiments, single-instruction multiple-thread (SIMT) techniques are used to support parallel execution of a large number of generally synchronized threads, using a common instruction unit configured to issue instructions to a set of cores.

(30) Each of the functional units **104** may connect to a URF **108** and a constant cache **110** via an interconnect network, for example, a hierarchical cross-bar with one or more read and/or write crossbars. The constant cache **110** may provide low-latency on-chip memory near the subpartition units of a multiprocessor **102**. The uniform register file **108** may include data registers assignable to

a different functional unit of the plurality of functional units **104** and/or different warps being executed by the multiprocessor **102**. The register file **108** provides temporary storage for functional units **104** on the multiprocessor **102**.

(31) The PPU **100** may support multiple address spaces including local, shared, and global address spaces to support data visibility for the threads. Additional read-only address spaces including textures may be supported. Each thread has its own per thread local or private memory which may be controlled by allocation of registers.

(32) Subpartition units **104** may request access to constant data stored in the constant cache **110** and/or stored in the constant memory **114** through the request coalescer **120**. Each subpartition unit **104** in the multiprocessor **102** is coupled the request coalescer **120** which is configured to access constant memory **114**. Each request coalescer **120** has read access to a constant cache of the corresponding multiprocessor **102** by issuing requests to the memory subsystem, such as by issuing requests to a general cache controller (GCC).

(33) When one or more of the subpartition units **104** require access to constant data stored in the constant cache **110**, the software initiates a thread with a “load” or “read” instruction. The load instruction may cause constant data to be read from the constant cache **110** into one or more uniform registers of the URF **108**, making it useable by all threads in the corresponding subpartition unit **104**. In at least one embodiment, the constant data may be written from the constant memory **114** into the constant cache **110** in the event of cache miss.

(34) The request coalescer **120** may coalesce multiple requests/instructions issued by one or more subpartition units **104** of a multiprocessor **102** to multicast constant data into one or more respective URFs **108**. In at least one embodiment, the request coalescer **120** may include a data structure (e.g., a content addressable memory (CAM) table) in which each entry of the request coalescer includes one or more requests/instructions to load the same constant data into respective URFs. The request coalescer **120** may coalesce multiple requests/instructions into the same entry based on an instruction identifier associated with the instructions. For example, the request coalescer may coalesce a first instruction and a second instruction with the same instruction identifier into the same entry of the request coalescer **120**. Matching instruction identifiers may indicate that the first instruction and the second instruction are the same instruction (e.g., the same instruction but associated with different warps and executing on different subpartition units **104**), and thus require the same constant data for execution. The request coalescer **120** may cause that same constant data to be multicast from the constant cache **110** into an URF **108A** associated with the first request and an URF **108B** associated with the second instruction.

(35) The request coalescer **120** may begin a process of multicasting constant data from the constant cache **110** into respective URFs **108** by selecting an eligible entry among multiple eligible entries. An eligible entry may be defined as an entry in the request coalescer **120** with a pending instruction to load constant data into URF(s) **108** that is currently available in the constant cache **110**. When eligible entry is selected, associated constant data can be pushed into WB buffers **109** of each subpartition unit **104** associated with the selected entry. The WB buffers **109** may temporarily hold constant data scheduled for writeback into associated URFs **108**. In at least one embodiment, the WB buffer **109** is a first-in, first-out (FIFO) queue. In some instances, constant data stored in the WB buffers **109** can be written into associated URFs **108** per subpartition **104** in request order.

(36) FIG. 2 illustrates a conceptual diagram **200** as instructions requesting constant data are coalesced and the constant data is multicast to associated uniform register files (URF(s)), in accordance with at least one embodiment of the present disclosure. An individual thread or a warp executing on a subpartition unit, such as subpartition unit **104** of FIG. 1, may request access to constant data from a constant cache **110** associated with a multiprocessor and load the constant data into one or more registers of a uniform register file (e.g., URF **108**) associated with the multiprocessor. In at least one embodiment, multiple requests may be received in a same clock cycle. For example, a request **202A**, a request **202B**, a request **202C**, and a request **202N** (referred

to generally as request(s) **202** herein) may be received in the same clock cycle. In some embodiments, each of the requests **202** may be associated with a warp that executes the same instruction simultaneously using SIMT architecture described with respect to FIG. 1. For example, the request **202A** may be associated with a warp executing an instruction on a subpartition unit **104A**, and the request **202B** may be associated with a different warp executing an instruction on a subpartition unit **104B**. One or more threads within the warp may request constant data from the constant cache to continue execution. In some instances, threads across multiple warps may request the same constant data. For example, the requests **202A** and **202B** may be associated with respective warps executing on respective subpartition units requesting the same constant data. In at least one embodiment, the requests **202A** and **202B** may be associated with a same instruction (e.g., the same instruction parameters).

(37) In at least one embodiment, the request **202** may be an instruction associated with a thread to be executed and may also be referred to as an “instruction **202**” herein. In at least one embodiment, the instruction **202** may include parameters that determine the operation to be performed such as operation code (Opcode), source operands, destination operands, and/or execution context. The Opcode, for example, may specify the operation or computation that the thread is instructed to perform, including arithmetic operations (addition, multiplication, etc.), logical operation (AND, OR, XOR, etc.), memory operations (load, store, etc.), control flow operations (branch, looping, etc.), and other instructions. The source operands are data on which the operation is to be performed or the address of the data, including registers, constants, data from global memory, etc. The destination operand may include the location where the result of the operation is to be stored, such as a register, a location in device memory, etc. The execution context may include information pertaining to a current state of the thread, including a program counter (PC), instruction pointer, register value, or the like.

(38) N to 1 arbiter **204** may arbitrate access into the coalescing pipeline **210** for coalescing of requests **202A** through **202N** and multicasting of associated constant data. The N to 1 arbiter **204** may be implemented using any suitable arbitration scheme, including, but not limited to, round-robin arbitration, static priority arbitration, rotational priority arbitration, fair share arbitration, dynamic priority arbitration, token passing arbitration, reservation-based arbitration, contention-based arbitration, and the like. In at least one embodiment, a Memory Input/Output Controller (MIOC) **201** may deliver instructions/requests to coalescing pipeline **210**. In at least one embodiment, the arbitration may be performed by MIOC **201**.

(39) At a tag lookup **212** stage of the coalescing pipeline **210**, a tag lookup may be performed to determine whether a requested constant is present in the constant cache **110**. An address included in the request **202** may be partitioned into a tag, an index, and an offset to identify the location of the requested constant in the constant cache **110**. In at least one embodiment, a general cache controller (GCC) **206** may perform the tag lookup by comparing the tag of the address with tags stored in the constant cache **110**, where each cache line in the constant cache **110** includes an associated tag. If the compare results in a match, it is a cache hit. If the compare does not result in a match, it is a cache miss. In at least one embodiment, the constant cache **110** is connected to the general cache controller **206** (GCC) through an interface. In the event of a miss on the constant cache **110**, a request may be issued to the GCC **206** to load data from the constant memory **114** into the constant cache **110** and the GCC **206** can manage the transfer of constant data from the constant memory to the constant cache **110**.

(40) Whether the tag lookup **212** stage results in a cache hit or a cache miss, the request **202** may be stored in a request coalescer **120** pending writeback of the constant data into URF(s) **108**. According to implementations described herein, writeback of constant data can refer to the process of loading data from constant cache **110** to one or more URFs **108**. The request coalescer **120** is a mechanism designed to merge or combine multiple requests **202** corresponding to the same constant data/same instruction such that warps issuing the same instruction are not serviced

individually. Specifically, the request coalescer **120** may be a specialized hardware structure with multiple entries that each correspond to a constant data and associated instructions/requests. Each entry of the request coalescer **120** may include instructions to writeback/load the constant data into associated URFs **108**. In at least one embodiment, each entry of the request coalescer **120** is associated with an instruction identifier and multiple instructions with the same instruction identifier can be coalesced into a same entry of the request coalescer **120** pending writeback.

(41) In an illustrative example, the request coalescer **120** may be a Content Addressable Memory (CAM). Each entry of the CAM may include a lookup portion and a payload portion. The lookup portion may include an instruction identifier to match incoming requests **202** against. The instruction identifier may include any combination of instruction parameters associated with the income requests **202**. For example, the instruction parameters may include cache line identifiers (e.g., a set identifier, a way identifier, etc.), a destination uniform register, a load size of a read operation of the request **202**, a scoreboard identifier, a byte offset associated with the cache line, an Opcode, a cache tag identifier, and the like. Generally, each entry of the CAM may include sufficient information to ensure it is functionally correct and/or acceptable to coalesce two or more requests for constant data. The payload may include an indication of the requests **202** requesting access to the same constant data, and therefore warps issuing the same instruction. In an embodiment where each request **202** is associated with a particular warp, the payload may include a bitmask indicating warps awaiting the same constant data for writeback into associated uniform registers, as illustrated below with respect to FIG. 3.

(42) Responsive to reception of a request **202** into the request coalescer **120**, the request coalescer **120** may compare the received request **202** against existing entries of the request coalescer **120**. As mentioned above, each entry of the request coalescer **120** may be associated with an instruction identifier such as one or more instruction parameters. If the instruction identifier of the request **202** matches an existing entry of the request coalescer **120**, the request **202** may be merged into the existing entry. If the instruction identifiers do not match, a new entry may be allocated and associated with the instruction identifier of the request **20** and the request **202** may be merged into the newly-allocated request coalescer **120** entry.

(43) Responsive to constant data being available in the constant cache **110**, the request coalescer **120** may begin the writeback process. To begin the writeback process, the request coalescer **120** selects an entry among entries eligible for writeback. An eligible entry is an entry within the request coalescer with a pending request **202** associated with constant data available to be read out from the constant cache **110**. In at least one embodiment, the request coalescer **120** may select from multiple eligible entries using an oldest-age first approach by comparing age fields associated with each eligible entry. The GCC **206** may facilitate loading constant data associated with the selected entry into associated writeback buffers **109**.

(44) The writeback buffer(s) **109** may temporarily hold constant data selected by the request coalescer **120** for writeback into associated URFs **108**. In at least one embodiment, a single writeback buffer **109** can be associated with a single subpartition unit and associated URF **108**, as illustrated above with respect to FIG. 1. In an illustrative example, a request **202A** and a request **202B** may be the same instruction associated (requesting the same constant data) with a subpartition **104A** and a subpartition **104B**, respectively. The request coalescer **120** can coalesce (e.g., merge) the request **202A** and **202B** into a given, same entry of the request coalescer **120**. Responsive to the request coalescer **120** selecting the given entry for writeback, the constant data may be pushed into the writeback buffer **109A** associated with the subpartition **104A** and a the writeback buffer **109B** associated with the subpartition **104B**. Constant data can be loaded from the writeback buffers **109A** and **109B** to URF **108A** and **108B**, respectively. Accordingly, constant data can be multicast from writeback buffers **109** to URFs **108** in a single clock cycle. In at least one embodiment, when a writeback buffer **109** become full, the writeback buffer **109** can backpressure the request coalescer **120** from unloading entries into the writeback buffer **109**.

(45) In some embodiments, coalesced requests associated with a same subpartition **104** may be serialized as they are written from a writeback buffer **109** to an associated URF **108**. For example, subpartition **104A** may include four warps with a constant cache load instruction that are coalesced into a single entry of the request coalescer **120** entry according to coalescing pipeline **210** described above. When the entry of the request coalescer **120** is selected for writeback, associated constant data may be written to a writeback buffer **109A** along with warp identifiers associated with warps of the request coalescer entry. Writeback of data from writeback buffer **109A** to an associated URF **108A** can be serialized for each coalesced warp across N cycles, where N is the number of coalesced warps in the writeback buffer **109A**. In at least one embodiment, writeback of warps can be serialized according to warp identifiers. For example, a first warp with a warp identifier '0' may be written back to URF **108A** followed by a second warp with a warp identifier '1.' In at least one embodiment, writeback of warps can be serialized according to an order of coalescing.

(46) In at least one embodiment, a given subpartition unit may include multiple URFs, associated with one or more warps executing on the given subpartition unit. For example, the given subpartition unit may include per-warp URFs such that each warp executing on the given subpartition unit utilizes a dedicated URF. In such an embodiment, the given subpartition unit may writeback data into such per-warp URFs in a single cycle. Generally, N-warps of a given subpartition unit may be merged into a set of M-URFs such that writeback serialization would require at least one cycle and at most N-cycles to writeback data.

(47) FIG. 3 illustrates an example request coalescer Content Addressable Memory (CAM) table **300**, in accordance with at least one embodiment of the present disclosure. In at least one embodiment, the request coalescer CAM **300** may correspond to the request coalescer **120** of FIG. 2. In at least one embodiment, the request coalescer CAM **300** may be a specialized memory for coalescing requests to write constant data from a constant cache/constant memory into associated register files, as described above. In at least one embodiment, control and management of the request coalescer CAM **300** may be facilitated through a hardware component, firmware, and/or one or more software layers.

(48) The request coalescer CAM **300** may include multiple entries **304A** through **304N** (referred to generally as "entry(ies) **304**" herein). Each entry may include a lookup and a payload. For example, the entry **304A** may include a lookup portion **302** and a payload portion **322**. The lookup portion **320** may include instruction identifiers to match against. The Lookup portion **320** may include, but is not limited to, any combination of a destination register **306**, a scoreboard identifier **308**, a cacheline ID **310**, a byte offset **311**, a load size **312**, an opcode **314**, and/or a tag **316**. The destination register **306** may be an identifier of a destination register associated with one or more URFs **108**. The scoreboard identifier **308** may include an identifier of a scoreboard to be released upon writeback. A scoreboard is a mechanism used for out-of-order instruction execution dependency tracking. The cacheline ID **310** may include a set identifier and a way identifier associated with the constant cache **110**. The byte offset **311** may be an identifier of a specific byte within the cacheline where the constant data is located. The load size **312** can indicate a size of associated constant data (e.g., 32 bits, 64 bits, 96 bits, etc.). The Opcode **314** may be an instruction code that represents a specific operation to be performed by the thread. The opcode **314** may include arithmetic operations (e.g., add, multiply, etc.), logical operations (e.g., AND, OR, etc.), memory operations (e.g., load, store, etc.), and control flow operations (e.g., branch, looping, etc.).

(49) The lookup portion **320** may include other various fields including information (e.g., metadata) not associated with an instruction identifier. In at least one embodiment, the lookup portion **320** may include a field to indicate whether the entry **304A** is valid. In at least one embodiment, the portion **302** may include to indicate whether the destination register **306** is valid. In at least one embodiment, the lookup portion **320** may include an indicator of whether constant data associated with the entry **304A** is available for writeback. In at least one embodiment, the lookup portion **320** may include a field that indicates whether an instruction/request in the entry

304A has a same instruction identifier as another instruction in the request coalescer CAM **300** associated with the same warp. In at least one embodiment, the lookup portion **320** may include a field that indicates an age of the entry **304A**.

(50) The payload portion **322** may include a bitmask identifying warps awaiting the same constant data associated with the entry **304A**. Bit position within the bitmask can represent an identifier associated with a warp. If the bit is set to '1', the corresponding warp is indicated in the entry as awaiting the constant data; if the bit is set to '0', the corresponding warp is not awaiting constant data. In an illustrative example, the bitmask "1000010000001000" may represent a payload **322** of the entry **304A** where the 0th, 5th, and 12th warp are awaiting the constant data associated with the entry **304A**. It is appreciated that the request coalescer CAM **300** is used herein by way of example, and not by way of limitation, noting that other associative memory solutions and implementations of a request coalescer may be utilized to coalesce requests. For example, a request coalescer may implement a hash table, a tree structure, or other associative memories.

(51) FIG. **3B** illustrates example pseudocode to coalesce a request into the example request coalescer CAM table **300**, in accordance with at least one embodiment of the present disclosure. The pseudocode is arranged to process a request through a series of if/else statements. The request is processed through a first if statement in which a for loop iterates through each entry of the CAM table **300**. If the request matching an existing entry and the exiting entry is able to coalesce (coalescible), the request is coalesced into the existing entry and the process ends. The request matches the existing entry if the existing entry is valid and the destination register, read size, scoreboard ID, byte offset, set ID, and way ID of the request matches that of the existing entry. The existing entry is coalescible if the existing entry is not locked, and warp IDs of the existing entry and the request do not match.

(52) Continuing onto FIG. **3C**, else if, the warp IDs match, then the request may not coalesce, and the pipeline is stalled.

(53) Else if, a new entry is needed in the CAM table **300**, then it is determined whether the CAM table **300** has an unused entry. If the CAM table **300** has an unused entry, then a new entry may be allocated to the unused entry. Once allocated, the set ID, the way ID, the destination uniform register, the read size, the scoreboard ID, and the byte offset associated with the new entry is set to corresponding values associated with the request. Additionally, the new entry is set to valid; the new entry is unlocked; a data ready indicator is set to true if the tag associated with the request hit; the warp mask of the new entry is set to include the warp associated with the request, and the new entry is set to the youngest entry in the CAM table **300**.

(54) Else, a new entry is needed but there is no unused entry in the CAM table **300**, and the pipeline is stalled.

(55) FIG. **4** illustrates a flowchart of a method **400** for multicasting data, in accordance with at least one embodiment of the present disclosure. Although method **400** is described in the context of a processing unit, the method **400** may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method **400** may be executed by a GPU (graphics processing unit), CPU (central processing unit), or any processor capable of multicasting data. Furthermore, persons of ordinary skill in the art will understand that any system that performs method **400** is within the scope and spirit of embodiments of the present invention.

(56) More illustrative information will now be set forth regarding various optional architectures and features with which the foregoing framework may be implemented, per the desires of the user. It should be strongly noted that the following information is set forth for illustrative purposes and should not be construed as limiting in any manner. Any of the following features may be optionally incorporated with or without the exclusion of other features described.

(57) Although shown in a particular sequence or order, unless otherwise specified, the order of the operations can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated operations can be performed in a different order, and some operations

can be performed in parallel. Additionally, one or more operations can be omitted in various embodiments. Thus, not all operations are required in every embodiment.

(58) At operation **402** of method **400**, processing logic, such as request coalescer **120** of FIG. **1**, may receive a first instruction to load a first data into a first uniform register file (URF) associated with a first subpartition unit, such as a subpartition unit **104A** of FIG. **1**.

(59) At operation **404** of method **400**, the processing logic may receive a second instruction to load the first data into a second URF associated with a second subpartition unit, such as a subpartition unit **104B** of FIG. **1**. In at least one embodiment, the first instruction and the second instruction are a same instruction. In at least one embodiment, the first instruction and the second instruction may be the same instruction in that they are the same instruction executed by different threads or warps rather than the same instance of a single instruction, as executed by a single thread. In at least one embodiment, instructions are received into the request coalescer according to an arbitration scheme, such as a round robin arbitration scheme.

(60) At operation **406** of method **400**, the processing logic may coalesce the first instruction and the second instruction into a first entry of the request coalescer based on instruction identifiers. The first entry is associated with the first data. In at least one embodiment, the coalesce the first instruction and the second instruction into the first entry of the request coalescer, the processing logic is to identity an instruction identifier associated with the first entry of the request coalescer. The processing logic is further to determine that the instruction identifier associated with the first entry matches respective instruction identifiers associated with the first instruction and the second instruction, and merge the first instruction and the second instruction into the first entry. In at least one embodiment, instruction identifiers can include at least one of a destination register number, a write scoreboard identifier, a load size, an opcode, a cache line identifier, a cache tag identifier, or any combination thereof.

(61) At operation **408** of method **400**, the processing logic may determine that the first data is available in a cache, such as constant cache **110** of FIG. **1**. In at least one embodiment, the cache is associated with the first partition unit and the second partition unit. In at least one embodiment, the constant cache is configured to stored constant data.

(62) At operation **410** of method **400**, responsive to the determination that the first data is available in the cache, the processing logic may multicast the first data to the first uniform register file and the second uniform register file. In at least one embodiment, responsive to the multicast of the first data from the cache to the first URF and the second URF, the processing logic may remove the first entry from the request coalescer. In at least one embodiment, the cache is a read only cache. In at least one embodiment, to multicast the data, the processing logic may write the data to a first writeback buffer and a second writeback buffer. The first writeback buffer can be associated with the first uniform register file, and the processing logic can write the first data to the first writeback buffer. The first writeback buffer can hold the first data until the first uniform register file is available for writeback. The second writeback buffer can be associated with the second uniform register file, and the processing logic can write the first data to the second writeback buffer. The second writeback buffer can hold the first data until the second uniform register file is available for writeback.

(63) Parallel Processing Architecture

(64) FIG. **5** illustrates a parallel processing unit (PPU) **500**, in accordance with an embodiment. In an embodiment, the PPU **500** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The PPU **500** is a latency hiding architecture designed to process many threads in parallel. A thread (e.g., a thread of execution) is an instantiation of a set of instructions configured to be executed by the PPU **500**. In an embodiment, the PPU **500** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device such as a liquid crystal display (LCD) device. In other embodiments, the PPU

500 may be utilized for performing general-purpose computations. While one exemplary parallel processor is provided herein for illustrative purposes, it should be strongly noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

(65) One or more PPU **500** may be configured to accelerate thousands of High Performance Computing (HPC), data center, and machine learning applications. The PPU **500** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image, and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and the like.

(66) As shown in FIG. 5, the PPU **500** includes an Input/Output (I/O) unit **505**, a front-end unit **515**, a scheduler unit **520**, a work distribution unit **525**, a hub **530**, a crossbar (Xbar) **570**, one or more processing clusters **550** (e.g., general processing clusters (GPCs), and one or more partition units **580**. The PPU **500** may be connected to a host processor or other PPUs **500** via one or more high-speed NVLink **510** interconnect. The PPU **500** may be connected to a host processor or other peripheral devices via an interconnect **502**. The PPU **500** may also be connected to a local memory comprising a number of memory devices **504**. In an embodiment, the local memory may comprise a number of dynamic random-access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device.

(67) The NVLink **510** interconnect enables systems to scale and include one or more PPUs **500** combined with one or more CPUs, supports cache coherence between the PPUs **500** and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **510** through the hub **530** to/from other units of the PPU **500** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). The NVLink **510** is described in more detail in conjunction with FIG. 5B.

(68) The I/O unit **505** is configured to transmit and receive communications (e.g., commands, data, etc.) from a host processor (not shown) over the interconnect **502**. The I/O unit **505** may communicate with the host processor directly via the interconnect **502** or through one or more intermediate devices such as a memory bridge. In an embodiment, the I/O unit **505** may communicate with one or more other processors, such as one or more the PPUs **500** via the interconnect **502**. In an embodiment, the I/O unit **505** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect **502** is a PCIe bus. In alternative embodiments, the I/O unit **505** may implement other types of well-known interfaces for communicating with external devices.

(69) The I/O unit **505** decodes packets received via the interconnect **502**. In an embodiment, the packets represent commands configured to cause the PPU **500** to perform various operations. The I/O unit **505** transmits the decoded commands to various other units of the PPU **500** as the commands may specify. For example, some commands may be transmitted to the front-end unit **515**. Other commands may be transmitted to the hub **530** or other units of the PPU **500** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **505** is configured to route communications between and among the various logical units of the PPU **500**.

(70) In an embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the PPU **500** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer is a region in a memory that is accessible (e.g., read/write) by both the host processor and the PPU **500**. For example, the I/O unit **505** may be configured to access the buffer in a system memory connected to the interconnect

502 via memory requests transmitted over the interconnect **502**. In an embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the PPU **500**. The front-end unit **515** receives pointers to one or more command streams. The front-end unit **515** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the PPU **500**.

(71) The front-end unit **515** is coupled to a scheduler unit **520** that configures the various processing clusters **550** to process tasks defined by the one or more streams. The scheduler unit **520** is configured to track state information related to the various tasks managed by the scheduler unit **520**. The state may indicate which processing cluster **550** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **520** manages the execution of a plurality of tasks on the one or more processing clusters **550**.

(72) The scheduler unit **520** is coupled to a work distribution unit **525** that is configured to dispatch tasks for execution on the processing clusters **550**. The work distribution unit **525** may track a number of scheduled tasks received from the scheduler unit **520**. In an embodiment, the work distribution unit **525** manages a pending task pool and an active task pool for each of the processing clusters **550**. The pending task pool may comprise a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular processing cluster **550**. The active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by the processing clusters **550**. As a processing cluster **550** finishes the execution of a task, that task is evicted from the active task pool for the processing cluster **550** and one of the other tasks from the pending task pool is selected and scheduled for execution on the processing cluster **550**. If an active task has been idle on the processing cluster **550**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the processing cluster **550** and returned to the pending task pool while another task in the pending task pool is selected and scheduled for execution on the processing cluster **550**.

(73) The work distribution unit **525** communicates with the one or more processing clusters **550** via XBar **570**. The XBar **570** is an interconnect network that couples many of the units of the PPU **500** to other units of the PPU **500**. For example, the XBar **570** may be configured to couple the work distribution unit **525** to a particular processing cluster **550**. Although not shown explicitly, one or more other units of the PPU **500** may also be connected to the XBar **570** via the hub **530**.

(74) The tasks are managed by the scheduler unit **520** and dispatched to a processing cluster **550** by the work distribution unit **525**. The processing cluster **550** is configured to process the task and generate results. The results may be consumed by other tasks within the processing cluster **550**, routed to a different processing cluster **550** via the XBar **570**, or stored in the memory **504**. The results can be written to the memory **504** via the partition units **580**, which implement a memory interface for reading and writing data to/from the memory **504**. The results can be transmitted to another PPU **500** or CPU via the NVLink **510**. In an embodiment, the PPU **500** includes a number U of partition units **580** that is equal to the number of separate and distinct memory devices **504** coupled to the PPU **500**. A partition unit **580** will be described in more detail below in conjunction with FIG. 6B.

(75) In an embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the PPU **500**. In an embodiment, multiple compute applications are simultaneously executed by the PPU **500** and the PPU **500** provides isolation, quality of service (QOS), and independent address spaces for the multiple compute applications. An application may generate instructions (e.g., API calls) that cause the driver kernel to generate one or more tasks for execution by the PPU **500**. The driver kernel outputs tasks to one or more streams being processed by the PPU **500**. Each task may comprise one or more groups of related threads, referred to herein as a warp. In an embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including

instructions to perform the task and that may exchange data through shared memory. Threads and cooperating threads are described in more detail in conjunction with FIG. 7A.

(76) FIG. 6A illustrates a processing cluster **550** of the PPU **500** of FIG. 5, in accordance with an embodiment. As shown in FIG. 6A, each processing cluster **550** includes a number of hardware units for processing tasks. In an embodiment, each processing cluster **550** includes a pipeline manager **610**, a pre-raster operations unit (PROP) **615**, a raster engine **625**, a work distribution crossbar (WDX) **680**, a memory management unit (MMU) **690**, and one or more Data Processing Clusters (DPCs) **620**. It will be appreciated that the processing cluster **550** of FIG. 6A may include other hardware units in lieu of or in addition to the units shown in FIG. 6A.

(77) In an embodiment, the operation of the processing cluster **550** is controlled by the pipeline manager **610**. The pipeline manager **610** manages the configuration of the one or more DPCs **620** for processing tasks allocated to the processing cluster **550**. In an embodiment, the pipeline manager **610** may configure at least one of the one or more DPCs **620** to implement at least a portion of a graphics rendering pipeline. For example, a DPC **620** may be configured to execute a vertex shader program on the programmable multiprocessor **640**. Multiprocessors **640** may generally include streaming multiprocessors, compute units, many integrated cores, and the like. The pipeline manager **610** may also be configured to route packets received from the work distribution unit **525** to the appropriate logical units within the processing cluster **550**. For example, some packets may be routed to fixed function hardware units in the PROP **615** and/or raster engine **625** while other packets may be routed to the DPCs **620** for processing by the primitive engine **635** or the multiprocessor **640**. In an embodiment, the pipeline manager **610** may configure at least one of the one or more DPCs **620** to implement a neural network model and/or a computing pipeline.

(78) The PROP unit **615** is configured to route data generated by the raster engine **625** and the DPCs **620** to a Raster Operations (ROP) unit, described in more detail in conjunction with FIG. 6B. The PROP unit **615** may also be configured to perform optimizations for color blending, organize pixel data, perform address translations, and the like.

(79) The raster engine **625** includes a number of fixed function hardware units configured to perform various raster operations. In an embodiment, the raster engine **625** includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, and a tile coalescing engine. The setup engine receives transformed vertices and generates plane equations associated with the geometric primitive defined by the vertices. The plane equations are transmitted to the coarse raster engine to generate coverage information (e.g., an x,y coverage mask for a tile) for the primitive. The output of the coarse raster engine is transmitted to the culling engine where fragments associated with the primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. Those fragments that survive clipping and culling may be passed to the fine raster engine to generate attributes for the pixel fragments based on the plane equations generated by the setup engine. The output of the raster engine **625** comprises fragments to be processed, for example, by a fragment shader implemented within a DPC **620**.

(80) Each DPC **620** included in the processing cluster **550** includes an M-Pipe Controller (MPC) **630**, a primitive engine **635**, and one or more Multiprocessors **640**. The MPC **630** controls the operation of the DPC **620**, routing packets received from the pipeline manager **610** to the appropriate units in the DPC **620**. For example, packets associated with a vertex may be routed to the primitive engine **635**, which is configured to fetch vertex attributes associated with the vertex from the memory **504**. In contrast, packets associated with a shader program may be transmitted to the multiprocessor **640**.

(81) In some embodiments, the multiprocessor **640** comprises a programmable multiprocessor, such as a programmable streaming multiprocessor, that is configured to process tasks represented by a number of threads. Each multiprocessor **640** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular warp concurrently. In an embodiment, the

multiprocessor **640** implements a SIMD (Single-Instruction, Multiple-Data) architecture where each thread in a warp is configured to process a different set of data based on the same set of instructions. All threads in the warp execute the same instructions. In another embodiment, the multiprocessor **640** implements a SIMT (Single-Instruction, Multiple Thread) architecture where each thread in a warp is configured to process a different set of data based on the same set of instructions, but where individual threads in the warp are allowed to diverge during execution. In an embodiment, a program counter, call stack, and execution state are maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call stack, and execution state are maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency. The multiprocessor **640** will be described in more detail below in conjunction with FIG. 7A.

(82) The MMU **690** provides an interface between the processing cluster **550** and the partition unit **580**. The MMU **690** may provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In an embodiment, the MMU **690** provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory **504**.

(83) FIG. 6B illustrates a memory partition unit **580** of the PPU **500** of FIG. 5, in accordance with an embodiment. As shown in FIG. 6B, the memory partition unit **580** includes a Raster Operations (ROP) unit **650**, a level two (L2) cache **660**, and a memory interface **670**. The memory interface **670** is coupled to the memory **504**. Memory interface **670** may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. In an embodiment, the PPU **500** incorporates U memory interfaces **670**, one memory interface **670** per pair of partition units **580**, where each pair of partition units **580** is connected to a corresponding memory device **504**. For example, PPU **500** may be connected to up to Y memory devices **504**, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random-access memory, or other types of persistent storage.

(84) In an embodiment, the memory interface **670** implements an HBM2 memory interface and Y equals half U. In an embodiment, the HBM2 memory stacks are located on the same physical package as the PPU **500**, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In an embodiment, each HBM2 stack includes four memory dies and Y equals 4, with HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

(85) In an embodiment, the memory **504** supports Single-Error Correcting Double-Error Detecting (SECCDED) Error Correction Code (ECC) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where PPUs **500** process very large datasets and/or run applications for extended periods.

(86) In an embodiment, the PPU **500** implements a multi-level memory hierarchy. In an embodiment, the memory partition unit **580** supports a unified memory to provide a single unified virtual address space for CPU and PPU **500** memory, enabling data sharing between virtual memory systems. In an embodiment the frequency of accesses by a PPU **500** to memory located on other processors is traced to ensure that memory pages are moved to the physical memory of the PPU **500** that is accessing the pages more frequently. In an embodiment, the NVLink **510** supports address translation services allowing the PPU **500** to directly access a CPU's page tables and providing full access to CPU memory by the PPU **500**.

(87) In an embodiment, copy engines transfer data between multiple PPUs **500** or between PPUs **500** and CPUs. The copy engines can generate page faults for addresses that are not mapped into the page tables. The memory partition unit **580** can then service the page faults, mapping the

addresses into the page table, after which the copy engine can perform the transfer. In a conventional system, memory is pinned (e.g., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing the available memory. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

(88) Data from the memory **504** or other system memory may be fetched by the memory partition unit **580** and stored in the L2 cache **660**, which is located on-chip and is shared between the various processing clusters **550**. As shown, each memory partition unit **580** includes a portion of the L2 cache **660** associated with a corresponding memory device **504**. Lower-level caches may then be implemented in various units within the processing clusters **550**. For example, each of the multiprocessors **640** may implement a level one (L1) cache. The L1 cache is private memory that is dedicated to a particular multiprocessor **640**. Data from the L2 cache **660** may be fetched and stored in each of the L1 caches for processing in the functional units of the Multiprocessors **640**. The L2 cache **660** is coupled to the memory interface **670** and the XBar **570**.

(89) The ROP unit **650** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and the like. The ROP unit **650** also implements depth testing in conjunction with the raster engine **625**, receiving a depth for a sample location associated with a pixel fragment from the culling engine of the raster engine **625**. The depth is tested against a corresponding depth in a depth buffer for a sample location associated with the fragment. If the fragment passes the depth test for the sample location, then the ROP unit **650** updates the depth buffer and transmits a result of the depth test to the raster engine **625**. It will be appreciated that the number of partition units **580** may be different than the number of processing clusters **550** and, therefore, each ROP unit **650** may be coupled to each of the processing clusters **550**. The ROP unit **650** tracks packets received from the different processing clusters **550** and determines which processing cluster **550** that a result generated by the ROP unit **650** is routed to through the Xbar **570**. Although the ROP unit **650** is included within the memory partition unit **580** in FIG. 6B, in other embodiment, the ROP unit **650** may be outside of the memory partition unit **580**. For example, the ROP unit **650** may reside in the processing cluster **550** or another unit.

(90) FIG. 7A illustrates the multiprocessor **640** of FIG. 6A, in accordance with an embodiment. As shown in FIG. 7A, the multiprocessor **640** includes an instruction cache **705**, one or more scheduler units **710**, a register file **720**, one or more processing cores **750**, one or more special function units (SFUs) **752**, one or more load/store units (LSUs) **754**, an interconnect network **780**, a shared memory/L1 cache **770**.

(91) As described above, the work distribution unit **525** dispatches tasks for execution on the processing clusters **550** of the PPU **500**. The tasks are allocated to a particular DPC **620** within a processing cluster **550** and, if the task is associated with a shader program, the task may be allocated to a multiprocessor **640**. The scheduler unit **710** receives the tasks from the work distribution unit **525** and manages instruction scheduling for one or more thread blocks assigned to the multiprocessor **640**. The scheduler unit **710** schedules thread blocks for execution as warps of parallel threads, where each thread block is allocated at least one warp. In an embodiment, each warp executes 32 threads. The scheduler unit **710** may manage a plurality of different thread blocks, allocating the warps to the different thread blocks and then dispatching instructions from the plurality of different cooperative groups to the various functional units (e.g., cores **750**, SFUs **752**, and LSUs **754**) during each clock cycle.

(92) Cooperative Groups is a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., the `syncthreads ( )` function). However,

programmers would often like to define warps at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

(93) Cooperative Groups enables programmers to define warps explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence.

Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(94) A dispatch unit **715** is configured to transmit instructions to one or more of the functional units. In the embodiment, the scheduler unit **710** includes two dispatch units **715** that enable two different instructions from the same warp to be dispatched during each clock cycle. In alternative embodiments, each scheduler unit **710** may include a single dispatch unit **715** or additional dispatch units **715**.

(95) Each multiprocessor **640** includes a register file **720** that provides a set of registers for the functional units of the multiprocessor **640**. In an embodiment, the register file **720** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of the register file **720**. In another embodiment, the register file **720** is divided between the different warps being executed by the multiprocessor **640**. The register file **720** provides temporary storage for operands connected to the data paths of the functional units.

(96) Each multiprocessor **640** comprises L processing cores **750**. In an embodiment, the multiprocessor **640** includes a large number (e.g., 128, etc.) of distinct processing cores **750**. Each core **750** may include a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating-point arithmetic logic unit and an integer arithmetic logic unit. In an embodiment, the floating-point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In an embodiment, the cores **750** include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(97) Tensor cores configured to perform matrix operations, and, in an embodiment, one or more tensor cores are included in the cores **750**. In particular, the tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In an embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D = A \times B + C$, where A, B, C, and D are 4×4 matrices.

(98) In an embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a $4 \times 4 \times 4$ matrix multiply. In practice, Tensor Cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16×16 size matrices spanning all 32 threads of the warp.

(99) Each multiprocessor **640** also comprises M SFUs **752** that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In an embodiment, the SFUs **752** may include a tree traversal unit configured to traverse a hierarchical tree data structure. In an embodiment, the SFUs **752** may include texture unit configured to perform texture map filtering operations. In an embodiment, the texture units are configured to load texture maps (e.g., a 2D

array of texels) from the memory **504** and sample the texture maps to produce sampled texture values for use in shader programs executed by the multiprocessor **640**. In an embodiment, the texture maps are stored in the shared memory/L1 cache **770**. The texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In an embodiment, each SM **540** includes two texture units.

(100) Each multiprocessor **640** also comprises N LSUs **754** that implement load and store operations between the shared memory/L1 cache **770** and the register file **720**. Each multiprocessor **640** includes an interconnect network **780** that connects each of the functional units to the register file **720** and the LSU **754** to the register file **720**, shared memory/L1 cache **770**. In an embodiment, the interconnect network **780** is a crossbar that can be configured to connect any of the functional units to any of the registers in the register file **720** and connect the LSUs **754** to the register file and memory locations in shared memory/L1 cache **770**.

(101) The shared memory/L1 cache **770** is an array of on-chip memory that allows for data storage and communication between the multiprocessor **640** and the primitive engine **635** and between threads in the multiprocessor **640**. In an embodiment, the shared memory/L1 cache **770** comprises 128 KB of storage capacity and is in the path from the multiprocessor **640** to the partition unit **580**. The shared memory/L1 cache **770** can be used to cache reads and writes. One or more of the shared memory/L1 cache **770**, L2 cache **660**, and memory **504** are backing stores.

(102) Combining data cache and shared memory functionality into a single memory block provides the best overall performance for both types of memory accesses. The capacity is usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory/L1 cache **770** enables the shared memory/L1 cache **770** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data.

(103) When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, the fixed function graphics processing units shown in FIG. 5, are bypassed, creating a much simpler programming model. In the general-purpose parallel computation configuration, the work distribution unit **525** assigns and distributes blocks of threads directly to the DPCs **620**. The threads in a block execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the multiprocessor **640** to execute the program and perform calculations, shared memory/L1 cache **770** to communicate between threads, and the LSU **754** to read and write global memory through the shared memory/L1 cache **770** and the memory partition unit **580**. When configured for general purpose parallel computation, the multiprocessor **640** can also write commands that the scheduler unit **520** can use to launch new work on the DPCs **620**.

(104) The PPU **500** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In an embodiment, the PPU **500** is embodied on a single semiconductor substrate. In another embodiment, the PPU **500** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional PPUs **500**, the memory **504**, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

(105) In an embodiment, the PPU **500** may be included on a graphics card that includes one or more memory devices **504**. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In yet another embodiment, the PPU **500** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

Exemplary Computing System

(106) Systems with multiple GPUs and CPUs are used in a variety of industries as developers

expose and leverage more parallelism in applications such as artificial intelligence computing. High-performance GPU-accelerated systems with tens to many thousands of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

(107) FIG. 7B is a conceptual diagram of a processing system **700** implemented using the PPU **500** of FIG. 5, in accordance with an embodiment. The exemplary system **765** may be configured to implement the method **400** shown in FIG. 4. The processing system **700** includes a CPU **730**, switch **712**, and multiple PPUs **500** each and respective memories **504**. The NVLink **510** provides high-speed communication links between each of the PPUs **500**. Although a particular number of NVLink **510** and interconnect **502** connections are illustrated in FIG. 7B, the number of connections to each PPU **500** and the CPU **730** may vary. The switch **712** interfaces between the interconnect **502** and the CPU **730**. The PPUs **500**, memories **504**, and NVLinks **510** may be situated on a single semiconductor platform to form a parallel processing module **725**. In an embodiment, the switch **712** supports two or more protocols to interface between various different connections and/or links.

(108) In another embodiment (not shown), the NVLink **510** provides one or more high-speed communication links between each of the PPUs **500** and the CPU **730** and the switch **712** interfaces between the interconnect **502** and each of the PPUs **500**. The PPUs **500**, memories **504**, and interconnect **502** may be situated on a single semiconductor platform to form a parallel processing module **725**. In yet another embodiment (not shown), the interconnect **502** provides one or more communication links between each of the PPUs **500** and the CPU **730** and the switch **712** interfaces between each of the PPUs **500** using the NVLink **510** to provide one or more high-speed communication links between the PPUs **500**. In another embodiment (not shown), the NVLink **510** provides one or more high-speed communication links between the PPUs **500** and the CPU **730** through the switch **712**. In yet another embodiment (not shown), the interconnect **502** provides one or more communication links between each of the PPUs **500** directly. One or more of the NVLink **510** high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink **510**.

(109) In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. It should be noted that the term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over utilizing a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the desires of the user. Alternately, the parallel processing module **725** may be implemented as a circuit board substrate and each of the PPUs **500** and/or memories **504** may be packaged devices. In an embodiment, the CPU **730**, switch **712**, and the parallel processing module **725** are situated on a single semiconductor platform.

(110) In an embodiment, the signaling rate of each NVLink **510** is 20 to 25 Gigabits/second and each PPU **500** includes six NVLink **510** interfaces (as shown in FIG. 7B, five NVLink **510** interfaces are included for each PPU **500**). Each NVLink **510** provides a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 500 Gigabytes/second. The NVLinks **510** can be used exclusively for PPU-to-PPU communication as shown in FIG. 7B, or some combination of PPU-to-PPU and PPU-to-CPU, when the CPU **730** also includes one or more NVLink **510** interfaces.

(111) In an embodiment, the NVLink **510** allows direct load/store/atomic access from the CPU **730** to each PPU's **500** memory **504**. In an embodiment, the NVLink **510** supports coherency operations, allowing data read from the memories **504** to be stored in the cache hierarchy of the CPU **730**, reducing cache access latency for the CPU **730**. In an embodiment, the NVLink **510**

includes support for Address Translation Services (ATS), allowing the PPU 500 to directly access page tables within the CPU 730. One or more of the NVLinks 510 may also be configured to operate in a low-power mode.

(112) FIG. 7C illustrates an exemplary system 765 in which the various architecture and/or functionality of the various previous embodiments may be implemented. The exemplary system 765 may be configured to implement the method 400 shown in FIG. 4.

(113) As shown, a system 765 is provided including at least one central processing unit 730 that is connected to a communication bus 775. The communication bus 775 may be implemented using any suitable protocol, such as PCI (Peripheral Component Interconnect), PCI-Express, AGP (Accelerated Graphics Port), HyperTransport, or any other bus or point-to-point communication protocol(s). The system 765 also includes a main memory 740. Control logic (software) and data are stored in the main memory 740 which may take the form of random-access memory (RAM).

(114) The system 765 also includes input devices 760, the parallel processing system 725, and display devices 745, e.g., a conventional CRT (cathode ray tube), LCD (liquid crystal display), LED (light emitting diode), plasma display or the like. User input may be received from the input devices 760, e.g., keyboard, mouse, touchpad, microphone, and the like. Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the system 765. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

(115) Further, the system 765 may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface 735 for communication purposes.

(116) The system 765 may also include a secondary storage (not shown). The secondary storage includes, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive reads from and/or writes to a removable storage unit in a well-known manner.

(117) Computer programs, or computer control logic algorithms, may be stored in the main memory 740 and/or the secondary storage. Such computer programs, when executed, enable the system 765 to perform various functions. The memory 740, the storage, and/or any other storage are possible examples of computer-readable media.

(118) The architecture and/or functionality of the various previous figures may be implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and/or any other desired system. For example, the system 765 may take the form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(119) While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary embodiments, but should be defined only in accordance with the following claims and their equivalents.

Graphics Processing Pipeline

(120) In an embodiment, the PPU 500 comprises a graphics processing unit (GPU). The PPU 500 is configured to receive commands that specify shader programs for processing graphics data.

Graphics data may be defined as a set of primitives such as points, lines, triangles, quads, triangle strips, and the like. Typically, a primitive includes data that specifies a number of vertices for the

primitive (e.g., in a model-space coordinate system) as well as attributes associated with each vertex of the primitive. The PPU **500** can be configured to process the graphics primitives to generate a frame buffer (e.g., pixel data for each of the pixels of the display).

(121) An application writes model data for a scene (e.g., a collection of vertices and attributes) to a memory such as a system memory or memory **504**. The model data defines each of the objects that may be visible on a display. The application then makes an API call to the driver kernel that requests the model data to be rendered and displayed. The driver kernel reads the model data and writes commands to the one or more streams to perform operations to process the model data. The commands may reference different shader programs to be implemented on the multiprocessors **640** of the PPU **500** including one or more of a vertex shader, hull shader, domain shader, geometry shader, and a pixel shader. For example, one or more of the multiprocessors **640** may be configured to execute a vertex shader program that processes a number of vertices defined by the model data. In an embodiment, the different multiprocessors **640** may be configured to execute different shader programs concurrently. For example, a first subset of multiprocessors **640** may be configured to execute a vertex shader program while a second subset of multiprocessors **640** may be configured to execute a pixel shader program. The first subset of multiprocessors **640** processes vertex data to produce processed vertex data and writes the processed vertex data to the L2 cache **460** and/or the memory **504**. After the processed vertex data is rasterized (e.g., transformed from three-dimensional data into two-dimensional data in screen space) to produce fragment data, the second subset of multiprocessors **640** executes a pixel shader to produce processed fragment data, which is then blended with other processed fragment data and written to the frame buffer in memory **504**. The vertex shader program and pixel shader program may execute concurrently, processing different data from the same scene in a pipelined fashion until all of the model data for the scene has been rendered to the frame buffer. Then, the contents of the frame buffer are transmitted to a display controller for display on a display device.

(122) FIG. **8** is a conceptual diagram of a graphics processing pipeline **800** implemented by the PPU **500** of FIG. **5**, in accordance with an embodiment. The graphics processing pipeline **800** is an abstract flow diagram of the processing steps implemented to generate 2D computer-generated images from 3D geometry data. As is well-known, pipeline architectures may perform long latency operations more efficiently by splitting up the operation into a plurality of stages, where the output of each stage is coupled to the input of the next successive stage. Thus, the graphics processing pipeline **800** receives input data **801** that is transmitted from one stage to the next stage of the graphics processing pipeline **800** to generate output data **802**. In an embodiment, the graphics processing pipeline **800** may represent a graphics processing pipeline defined by the OpenGL® API. As an option, the graphics processing pipeline **800** may be implemented in the context of the functionality and architecture of the previous Figures and/or any subsequent Figure(s).

(123) As shown in FIG. **8**, the graphics processing pipeline **800** comprises a pipeline architecture that includes a number of stages. The stages include, but are not limited to, a data assembly stage **810**, a vertex shading stage **820**, a primitive assembly stage **830**, a geometry shading stage **840**, a viewport scale, cull, and clip (VSCC) stage **850**, a rasterization stage **860**, a fragment shading stage **870**, and a raster operations stage **880**. In an embodiment, the input data **801** comprises commands that configure the processing units to implement the stages of the graphics processing pipeline **800** and geometric primitives (e.g., points, lines, triangles, quads, triangle strips or fans, etc.) to be processed by the stages. The output data **802** may comprise pixel data (e.g., color data) that is copied into a frame buffer or other type of surface data structure in a memory.

(124) The data assembly stage **810** receives the input data **801** that specifies vertex data for high-order surfaces, primitives, or the like. The data assembly stage **810** collects the vertex data in a temporary storage or queue, such as by receiving a command from the host processor that includes a pointer to a buffer in memory and reading the vertex data from the buffer. The vertex data is then transmitted to the vertex shading stage **820** for processing.

(125) The vertex shading stage **820** processes vertex data by performing a set of operations (e.g., a vertex shader or a program) once for each of the vertices. Vertices may be, e.g., specified as a 4-coordinate vector (e.g., $\langle x, y, z, w \rangle$) associated with one or more vertex attributes (e.g., color, texture coordinates, surface normal, etc.). The vertex shading stage **820** may manipulate individual vertex attributes such as position, color, texture coordinates, and the like. In other words, the vertex shading stage **820** performs operations on the vertex coordinates or other vertex attributes associated with a vertex. Such operations commonly including lighting operations (e.g., modifying color attributes for a vertex) and transformation operations (e.g., modifying the coordinate space for a vertex). For example, vertices may be specified using coordinates in an object-coordinate space, which are transformed by multiplying the coordinates by a matrix that translates the coordinates from the object-coordinate space into a world space or a normalized-device-coordinate (NCD) space. The vertex shading stage **820** generates transformed vertex data that is transmitted to the primitive assembly stage **830**.

(126) The primitive assembly stage **830** collects vertices output by the vertex shading stage **820** and groups the vertices into geometric primitives for processing by the geometry shading stage **840**. For example, the primitive assembly stage **830** may be configured to group every three consecutive vertices as a geometric primitive (e.g., a triangle) for transmission to the geometry shading stage **840**. In some embodiments, specific vertices may be reused for consecutive geometric primitives (e.g., two consecutive triangles in a triangle strip may share two vertices). The primitive assembly stage **830** transmits geometric primitives (e.g., a collection of associated vertices) to the geometry shading stage **840**.

(127) The geometry shading stage **840** processes geometric primitives by performing a set of operations (e.g., a geometry shader or program) on the geometric primitives. Tessellation operations may generate one or more geometric primitives from each geometric primitive. In other words, the geometry shading stage **840** may subdivide each geometric primitive into a finer mesh of two or more geometric primitives for processing by the rest of the graphics processing pipeline **800**. The geometry shading stage **840** transmits geometric primitives to the viewport SCC stage **850**.

(128) In an embodiment, the graphics processing pipeline **800** may operate within a multiprocessor and the vertex shading stage **820**, the primitive assembly stage **830**, the geometry shading stage **840**, the fragment shading stage **870**, and/or hardware/software associated therewith, may sequentially perform processing operations. Once the sequential processing operations are complete, in an embodiment, the viewport SCC stage **850** may utilize the data. In an embodiment, primitive data processed by one or more of the stages in the graphics processing pipeline **800** may be written to a cache (e.g., L1 cache, a vertex cache, etc.). In this case, in an embodiment, the viewport SCC stage **850** may access the data in the cache. In an embodiment, the viewport SCC stage **850** and the rasterization stage **860** are implemented as fixed function circuitry.

(129) The viewport SCC stage **850** performs viewport scaling, culling, and clipping of the geometric primitives. Each surface being rendered to is associated with an abstract camera position. The camera position represents a location of a viewer looking at the scene and defines a viewing frustum that encloses the objects of the scene. The viewing frustum may include a viewing plane, a rear plane, and four clipping planes. Any geometric primitive entirely outside of the viewing frustum may be culled (e.g., discarded) because the geometric primitive will not contribute to the final rendered scene. Any geometric primitive that is partially inside the viewing frustum and partially outside the viewing frustum may be clipped (e.g., transformed into a new geometric primitive that is enclosed within the viewing frustum). Furthermore, geometric primitives may each be scaled based on a depth of the viewing frustum. All potentially visible geometric primitives are then transmitted to the rasterization stage **860**.

(130) The rasterization stage **860** converts the 3D geometric primitives into 2D fragments (e.g., capable of being utilized for display, etc.). The rasterization stage **860** may be configured to utilize

the vertices of the geometric primitives to setup a set of plane equations from which various attributes can be interpolated. The rasterization stage **860** may also compute a coverage mask for a plurality of pixels that indicates whether one or more sample locations for the pixel intercept the geometric primitive. In an embodiment, z-testing may also be performed to determine if the geometric primitive is occluded by other geometric primitives that have already been rasterized. The rasterization stage **860** generates fragment data (e.g., interpolated vertex attributes associated with a particular sample location for each covered pixel) that are transmitted to the fragment shading stage **870**.

(131) The fragment shading stage **870** processes fragment data by performing a set of operations (e.g., a fragment shader or a program) on each of the fragments. The fragment shading stage **870** may generate pixel data (e.g., color values) for the fragment such as by performing lighting operations or sampling texture maps using interpolated texture coordinates for the fragment. The fragment shading stage **870** generates pixel data that is transmitted to the raster operations stage **880**.

(132) The raster operations stage **880** may perform various operations on the pixel data such as performing alpha tests, stencil tests, and blending the pixel data with other pixel data corresponding to other fragments associated with the pixel. When the raster operations stage **880** has finished processing the pixel data (e.g., the output data **802**), the pixel data may be written to a render target such as a frame buffer, a color buffer, or the like.

(133) It will be appreciated that one or more additional stages may be included in the graphics processing pipeline **800** in addition to or in lieu of one or more of the stages described above. Various implementations of the abstract graphics processing pipeline may implement different stages. Furthermore, one or more of the stages described above may be excluded from the graphics processing pipeline in some embodiments (such as the geometry shading stage **840**).

(134) Other types of graphics processing pipelines are contemplated as being within the scope of the present disclosure. Furthermore, any of the stages of the graphics processing pipeline **800** may be implemented by one or more dedicated hardware units within a graphics processor such as PPU **500**. Other stages of the graphics processing pipeline **800** may be implemented by programmable hardware units such as the multiprocessors **640** of the PPU **500**.

(135) The graphics processing pipeline **800** may be implemented via an application executed by a host processor, such as a CPU. In an embodiment, a device driver may implement an application programming interface (API) that defines various functions that can be utilized by an application in order to generate graphical data for display. The device driver is a software program that includes a plurality of instructions that control the operation of the PPU **500**. The API provides an abstraction for a programmer that lets a programmer utilize specialized graphics hardware, such as the PPU **500**, to generate the graphical data without requiring the programmer to utilize the specific instruction set for the PPU **500**. The application may include an API call that is routed to the device driver for the PPU **500**. The device driver interprets the API call and performs various operations to respond to the API call. In some instances, the device driver may perform operations by executing instructions on the CPU. In other instances, the device driver may perform operations, at least in part, by launching operations on the PPU **500** utilizing an input/output interface between the CPU and the PPU **500**. In an embodiment, the device driver is configured to implement the graphics processing pipeline **800** utilizing the hardware of the PPU **500**.

(136) Various programs may be executed within the PPU **500** in order to implement the various stages of the graphics processing pipeline **800**. For example, the device driver may launch a kernel on the PPU **500** to perform the vertex shading stage **820** on one multiprocessors **640** (or multiple multiprocessors **640**). The device driver (or the initial kernel executed by the PPU **500**) may also launch other kernels on the PPU **500** to perform other stages of the graphics processing pipeline **800**, such as the geometry shading stage **840** and the fragment shading stage **870**. In addition, some of the stages of the graphics processing pipeline **800** may be implemented on fixed unit hardware

such as a rasterizer or a data assembler implemented within the PPU **500**. It will be appreciated that results from one kernel may be processed by one or more intervening fixed function hardware units before being processed by a subsequent kernel on an multiprocessors **640**.

Machine Learning

(137) Deep neural networks (DNNs) developed on processors, such as the PPU **500** have been used for diverse use cases, from self-driving cars to faster drug development, from automatic image captioning in online image databases to smart real-time language translation in video chat applications. Deep learning is a technique that models the neural learning process of the human brain, continually learning, continually getting smarter, and delivering more accurate results more quickly over time. A child is initially taught by an adult to correctly identify and classify various shapes, eventually being able to identify shapes without any coaching. Similarly, a deep learning or neural learning system needs to be trained in object recognition and classification for it get smarter and more efficient at identifying basic objects, occluded objects, etc., while also assigning context to objects.

(138) At the simplest level, neurons in the human brain look at various inputs that are received, importance levels are assigned to each of these inputs, and output is passed on to other neurons to act upon. An artificial neuron or perceptron is the most basic model of a neural network. In one example, a perceptron may receive one or more inputs that represent various features of an object that the perceptron is being trained to recognize and classify, and each of these features is assigned a certain weight based on the importance of that feature in defining the shape of an object.

(139) A deep neural network (DNN) model includes multiple layers of many connected nodes (e.g., perceptrons, Boltzmann machines, radial basis functions, convolutional layers, etc.) that can be trained with enormous amounts of input data to quickly solve complex problems with high accuracy. In one example, a first layer of the DNN model breaks down an input image of an automobile into various sections and looks for basic patterns such as lines and angles. The second layer assembles the lines to look for higher level patterns such as wheels, windshields, and mirrors. The next layer identifies the type of vehicle, and the final few layers generate a label for the input image, identifying the model of a specific automobile brand.

(140) Once the DNN is trained, the DNN can be deployed and used to identify and classify objects or patterns in a process known as inference. Examples of inference (the process through which a DNN extracts useful information from a given input) include identifying handwritten numbers on checks deposited into ATM machines, identifying images of friends in photos, delivering movie recommendations to over fifty million users, identifying and classifying different types of automobiles, pedestrians, and road hazards in driverless cars, or translating human speech in real-time.

(141) During training, data flows through the DNN in a forward propagation phase until a prediction is produced that indicates a label corresponding to the input. If the neural network does not correctly label the input, then errors between the correct label and the predicted label are analyzed, and the weights are adjusted for each feature during a backward propagation phase until the DNN correctly labels the input and other inputs in a training dataset. Training complex neural networks requires massive amounts of parallel computing performance, including floating-point multiplications and additions that are supported by the PPU **500**. Inferencing is less compute-intensive than training, being a latency-sensitive process where a trained neural network is applied to new inputs it has not seen before to classify images, translate speech, and generally infer new information.

(142) Neural networks rely heavily on matrix math operations, and complex multi-layered networks require tremendous amounts of floating-point performance and bandwidth for both efficiency and speed. With thousands of processing cores, optimized for matrix math operations, and delivering tens to hundreds of TFLOPS of performance, the PPU **500** is a computing platform

capable of delivering performance required for deep neural network-based artificial intelligence and machine learning applications.

Claims

1. A parallel processing unit (PPU) comprising: a plurality of subpartition units configured to execute threads; a cache coupled to the plurality of subpartition units; and a request coalescer coupled to the plurality of subpartition units and the cache, wherein the request coalescer is to: receive a first instruction to load a first data into a first register file associated with a first subpartition unit of the plurality of subpartition units; receive a second instruction to load the first data into a second register file associated with a second subpartition unit of the plurality of subpartition units; coalesce the first instruction and the second instruction into a first entry of the request coalescer based on instruction identifiers, wherein the first entry is associated with the first data; determine that the first data is available in the cache; and responsive to the determination that the first data is available in the cache, multicast the first data from the cache to the first register file and the second register file.
2. The PPU of claim 1, further comprising: a first writeback buffer associated with the first register file, wherein the request coalescer is to write the first data to the first writeback buffer, wherein the first writeback buffer holds the first data until the first register file is available for writeback; and a second writeback buffer associated with the second register file, wherein the request coalescer is to write the first data to the first second writeback buffer, wherein the second writeback buffer holds the first data until the second register file is available for writeback.
3. The PPU of claim 2, wherein, to multicast the first data from the cache to the first register file and the second register file, the request coalescer is to: write the first data to the first writeback buffer and the second writeback buffer.
4. The PPU of claim 1, wherein the first instruction and the second instruction are a same instruction as executed by respective threads.
5. The PPU of claim 1, wherein the cache is configured to store constant data.
6. The PPU of claim 1, wherein, to coalesce the first instruction and the second instruction into the first entry of the request coalescer based on instruction identifiers, the request coalescer is to: identify an instruction identifier associated with the first entry of the request coalescer; determine that the instruction identifier associated with the first entry matches respective instruction identifiers associated with the first instruction and the second instruction; and merge the first instruction and the second instruction into the first entry.
7. The PPU of claim 6, wherein the instruction identifier comprises at least one of a destination register number, a write scoreboard identifier, a load size, an opcode, a cache line identifier, a program counter, or a cache tag identifier.
8. The PPU of claim 1, wherein the request coalescer is further to: responsive to the multicasting of the first data from the cache to the first register file and the second register file, remove the first entry from the request coalescer.
9. The PPU of claim 1, wherein instructions are received into the request coalescer according to an arbitration scheme.
10. The PPU of claim 1, wherein the cache is a read-only cache.
11. A parallel processor comprising: a plurality of multiprocessors, each having a cache and a request coalescer, wherein the request coalescer is to: receive a first instruction to load a first data into a first register file; receive a second instruction to load the first data into a second register file; coalesce the first instruction and the second instruction into a first entry of the request coalescer based on instruction identifiers, wherein the first entry is associated with the first data; determine that the first data is available in the cache; and responsive to the determination that the first data is available in the cache, multicast the first data from the cache to the first register file and the second

register file.

12. The parallel processor of claim 11, further comprising: a first writeback buffer associated with the first register file, wherein the request coalescer is to write the first data to the first writeback buffer, wherein the first writeback buffer holds the first data until the first register file is available for writeback; and a second writeback buffer associated with the second register file, wherein the request coalescer is to write the first data to the second writeback buffer, wherein the second writeback buffer holds the first data until the second register file is available for writeback.

13. The parallel processor of claim 12, wherein, to multicast the first data from the cache to the first register file and the second register file, the request coalescer is to: write the first data to the first writeback buffer and the second writeback buffer.

14. The parallel processor of claim 11, wherein each of the plurality of multiprocessors further comprises a plurality of subpartition units, wherein the first instruction is associated with a first subpartition unit of the plurality of subpartition units and the second instruction is associated with a second subpartition unit of the plurality of subpartition units.

15. The parallel processor of claim 11, wherein the first instruction and the second instruction are a same instruction as executed by respective threads.

16. The parallel processor of claim 11, wherein the cache is configured to store constant data.

17. The parallel processor of claim 11, wherein, to coalesce the first instruction and the second instruction into the first entry of the request coalescer, the request coalescer is to: identify an instruction identifier associated with the first entry of the request coalescer; determine that the instruction identifier associated with the first entry matches respective instruction identifiers associated with the first instruction and the second instruction; and merge the first instruction and the second instruction into the first entry.

18. The parallel processor of claim 17, wherein the instruction identifier comprises at least one of a destination register number, a write scoreboard identifier, a load size, an opcode, a cache line identifier, a program counter, or a cache tag identifier.

19. The parallel processor of claim 11, wherein the request coalescer is further to: responsive to the multicasting of the first data from the cache to the first register file and the second register file, remove the first entry from the request coalescer.

20. A method comprising: receiving a first instruction to load a first data into a first register file associated with a first subpartition unit of a plurality of subpartition units; receiving a second instruction to load the first data into a second register file associated with a second subpartition unit of the plurality of subpartition units; coalescing the first instruction and the second instruction into a first entry of a request coalescer based on instruction identifiers, wherein the first entry is associated with the first data; determining that the first data is available in a cache associated with the plurality of subpartition units; and responsive to determining that the first data is available in the cache, multicasting the first data from the cache to the first register file and the second register file.
