

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417180
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Benisty; Shay

Cache write overlap handling

Abstract

Embodiments of the present disclosure generally relate to a target device handling overlap write commands. In one embodiment, a target device includes a non-volatile memory and a controller coupled to the non-volatile memory. The controller includes a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer. The controller is configured to receive a new write command, classify the new write command, and write data associated with the new write command to one of the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer. Once the overlap accumulated buffer becomes available, the controller first flushes to the non-volatile memory the data in the random accumulated buffer and the sequential accumulated buffer that was received prior in sequence to the data in the overlap accumulated buffer. The controller then flushes the available overlap accumulated buffer, ensuring that new write commands override prior write commands.

Inventors:	Benisty; Shay (Beer Sheva, IL)
Applicant:	SANDISK TECHNOLOGIES, INC. (Milpitas, CA)
Family ID:	1000008819616
Assignee:	Sandisk Technologies, Inc. (Milpitas, CA)
Appl. No.:	18/515932
Filed:	November 21, 2023

Prior Publication Data

Document Identifier	Publication Date
US 20240086326 A1	Mar. 14, 2024

Related U.S. Application Data

division parent-doc US 17860263 20220708 US 11868257 child-doc US 18515932
division parent-doc US 16925191 20200709 US 11386006 20220712 child-doc US 17860263

Publication Classification

Int. Cl.: G06F12/08 (20160101); G06F3/06 (20060101); G06F12/0804 (20160101);
G06F12/0815 (20160101); G06F12/0855 (20160101); G06F12/0868 (20160101);
G06F12/0871 (20160101)

U.S. Cl.:

CPC G06F12/0815 (20130101); G06F3/061 (20130101); G06F3/0656 (20130101);
G06F3/0679 (20130101); G06F12/0804 (20130101); G06F12/0855 (20130101);
G06F12/0868 (20130101); G06F12/0871 (20130101); G06F2212/1024 (20130101);
G06F2212/1032 (20130101); G06F2212/1056 (20130101); G06F2212/214 (20130101);
G06F2212/312 (20130101); G06F2212/7203 (20130101)

Field of Classification Search

CPC: G06F (12/0815); G06F (12/0804); G06F (12/0855); G06F (12/0868); G06F (12/0871);
G06F (3/061); G06F (3/0656); G06F (3/0679); G06F (2212/1032); G06F (2212/1056);
G06F (2212/214); G06F (2212/312); G06F (2212/7203)

USPC: 711/143; 711/150; 711/151

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
10733100	12/2019	Benisty	N/A	N/A
2010/0281219	12/2009	Lippert	711/E12.001	G06F 12/0864
2011/0099323	12/2010	Syu	N/A	N/A
2012/0095966	12/2011	Yogev	N/A	N/A
2012/0198174	12/2011	Nellans et al.	N/A	N/A
2015/0026411	12/2014	Lippert et al.	N/A	N/A
2015/0309933	12/2014	Nellans et al.	N/A	N/A
2016/0110102	12/2015	Lee	711/103	G06F 13/1694
2017/0192889	12/2016	Sato et al.	N/A	N/A
2018/0018101	12/2017	Benisty et al.	N/A	N/A
2018/0039424	12/2017	Cui et al.	N/A	N/A
2020/0294616	12/2019	Helmich et al.	N/A	N/A
2020/0341897	12/2019	Benisty	N/A	N/A
2020/0371933	12/2019	Benisty	N/A	N/A

Primary Examiner: Nguyen; Than

Attorney, Agent or Firm: Patterson + Sheridan, LLP

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS (1) This application is a divisional application of U.S. patent application Ser. No. 17/860,263, filed on Jul. 8, 2022, which application is a divisional application of U.S. patent application Ser. No. 16/925,191, filed on Jul. 9, 2020, now U.S. Pat. No. 11,386,006, which application is a divisional application of U.S. patent application Ser. No. 16/003,914, filed on Jun. 8, 2018, now U.S. Pat. No. 10,733,100, which application claims benefit of U.S. Provisional Patent Application Ser. No. 62/518,260, filed Jun. 12, 2017, both of which are herein incorporated by reference.

BACKGROUND OF THE DISCLOSURE

Field of the Disclosure

(1) Embodiments of the present disclosure generally relate to a target device for handling overlap write scenarios, such as a solid state drive (SSD).

Description of the Related Art

(2) In order to increase write performance of a target device, such as a data storage device, the data storage device usually caches write data from a host device in a temporal buffer residing in either SRAM or DRAM before writing the data to non-volatile memory (NVM). The data storage device posts a completion message to the host even before writing the data to NVM.

(3) When supporting write caching mechanisms, a cache coherency logic is implemented in data storage devices in order to avoid coherency issues. For instance, when a data storage device receives two write commands that have a logical block address (LBA) overlap between them, the cache coherency logic does not allow the data associated with the first command received to override the data associated with the second command. When a data storage device detects cache coherency issues, the cache coherency logic either breaks the data pipeline or stops receiving any new host commands until the coherency issues are resolved by flushing all write cache data to NVM before receiving any new write commands. Thus, current solutions to cache coherency issues adversely add extra complexity and degrade performance of the data storage device.

(4) Therefore, there is a need for an improved data storage device.

SUMMARY OF THE DISCLOSURE

(5) Embodiments of the present disclosure generally relate to a target device handling overlap write commands. In one embodiment, a target device includes a non-volatile memory and a controller coupled to the non-volatile memory. The controller includes a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer. The controller is configured to receive a new write command, classify the new write command, and write data associated with the new write command to one of the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer. Once the overlap accumulated buffer becomes available, the controller first flushes to the non-volatile memory the data in the random accumulated buffer and the sequential accumulated buffer that was received prior in sequence to the data in the overlap accumulated buffer. The controller then flushes the available overlap accumulated buffer, ensuring that new write commands override prior write commands.

(6) In one embodiment, a data storage device comprises a non-volatile memory, and a controller coupled to the non-volatile memory. The controller comprises a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer. The controller is configured to receive a new write command, classify the new write command, and write data associated with the new write command to one of the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer based on the classification of the new write command.

(7) In another embodiment, a storage system comprises a host device, and a data storage device coupled to the host device. The data storage device comprises a non-volatile memory, and a

controller coupled to the non-volatile memory. The controller is configured to receive a new write command, push data associated to the new write command to one of a random stream, a sequential stream, or an overlap stream, and post a write command completion message to the host after the data associated with the new write command is pushed to one of the random stream, the sequential stream, or the overlap stream.

(8) In one embodiment, a method of handling write overlap data streams comprises receiving data to a controller in a sequence, detecting a cache collision of the data, and classifying the data as overlap data, random data, or sequential data. The method further comprises writing random data to a random buffer, writing sequential data to a sequential buffer, and writing overlap data to an overlap buffer, and emptying the overlap data buffer. Emptying the overlap data buffer comprises flushing the random data and the sequential data that were received prior to the overlap data in the sequence to a non-volatile memory, and flushing the overlap data to the non-volatile memory.

(9) In another embodiment, a method comprises detecting a first accumulated buffer of a controller is available to empty to a non-volatile memory, determining the first accumulated buffer is associated with an overlap stream, and delaying the emptying of the first accumulated buffer to the non-volatile memory. The method further comprises flushing a second accumulated buffer associated with a random stream, flushing a third accumulated buffer associated with a sequential stream, and emptying the first accumulated buffer to the non-volatile memory.

(10) In another embodiment, a data storage device comprises a non-volatile memory, means for detecting a cache collision, and means for ensuring that new write commands override prior write commands. The means for ensuring that new write commands override prior write commands comprises a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

(1) So that the manner in which the above recited features of the present disclosure can be understood in detail, a more particular description of the disclosure, briefly summarized above, may be had by reference to embodiments, some of which are illustrated in the appended drawings. It is to be noted, however, that the appended drawings illustrate only typical embodiments of this disclosure and are therefore not to be considered limiting of its scope, for the disclosure may admit to other equally effective embodiments.

(2) FIG. 1 is a schematic illustration of a system including a host and a target device, according to one embodiment.

(3) FIG. 2 is a method of a target device receiving data from a host and pushing the data to a cache, according to one embodiment.

(4) FIG. 3 is a method of writing data from a cache to non-volatile memory, according to one embodiment.

(5) To facilitate understanding, identical reference numerals have been used, where possible, to designate identical elements that are common to the figures. It is contemplated that elements disclosed in one embodiment may be beneficially utilized on other embodiments without specific recitation.

DETAILED DESCRIPTION

(6) In the following, reference is made to embodiments of the disclosure. However, it should be understood that the disclosure is not limited to specific described embodiments. Instead, any combination of the following features and elements, whether related to different embodiments or not, is contemplated to implement and practice the disclosure. Furthermore, although embodiments of the disclosure may achieve advantages over other possible solutions and/or over the prior art,

whether or not a particular advantage is achieved by a given embodiment is not limiting of the disclosure. Thus, the following aspects, features, embodiments and advantages are merely illustrative and are not considered elements or limitations of the appended claims except where explicitly recited in a claim(s). Likewise, reference to “the disclosure” shall not be construed as a generalization of any inventive subject matter disclosed herein and shall not be considered to be an element or limitation of the appended claims except where explicitly recited in a claim(s).

(7) Embodiments of the present disclosure generally relate to a target device handling overlap write commands. In one embodiment, a target device includes a non-volatile memory and a controller coupled to the non-volatile memory. The controller includes a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer. The controller is configured to receive a new write command, classify the new write command, and write data associated with the new write command to one of the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer. Once the overlap accumulated buffer becomes available, the controller first flushes to the non-volatile memory the data in the random accumulated buffer and the sequential accumulated buffer that was received prior in sequence to the data in the overlap accumulated buffer. The controller then flushes the available overlap accumulated buffer, ensuring that new write commands override prior write commands.

(8) FIG. 1 is a schematic illustration of a storage system **100**, according to one embodiment. The system **100** includes an initiator or host **150** and one or more target devices, such as a data storage device **102**, coupled to the host **150**. The host **150** utilizes a non-volatile memory (NVM) **106** included in the data storage device **102** to store and retrieve data. The data storage device **102** may be a removable mass storage device, such as, but not limited to, a handheld, removable memory device, such as a memory card (e.g., a Secure Digital (SD) card, a micro Secure Digital (micro-SD) card, or a MultiMedia Card (MMC)), or a universal serial bus (USB) device. The data storage device **102** may take the form of an embedded mass storage device, such as an eSD/eMMC embedded flash drive, embedded in host **150**.

(9) The host **150** may include a wide range of devices, such as computer servers, network attached storage (NAS) units, desktop computers, notebook (i.e., laptop) computers, tablet computers (i.e., “smart” pad), set-top boxes, telephone handsets (i.e., “smart” phones), televisions, cameras, display devices, digital media players, video gaming consoles, video streaming devices, and automotive applications (i.e., mapping, autonomous driving). In certain embodiments, the host **150** includes any device having a processing unit or any form of hardware capable of processing data, including a general purpose processing unit (such as a central processing unit (CPU)), dedicated hardware (such as an application specific integrated circuit (ASIC)), configurable hardware such as a field programmable gate array (FPGA), or any other form of processing unit configured by software instructions, microcode, or firmware.

(10) The host **150** interacts with the data storage device **102** through a host interface **101**. In certain embodiments, the storage system **100** operates following the non-volatile memory express (NVMe) protocol. In other embodiments, storage system **100** operates following Universal Flash Storage (UFS), serial advanced technology attachment (SATA), serially attached SCSI (SAS), advanced technology attachment (ATA), parallel-ATA (PATA), Fibre Channel Arbitrated Loop (FCAL), small computer system interface (SCSI), peripheral component interconnect (PCI), PCI-express (PCIe), and other suitable protocols.

(11) The NVM **106** of the data storage device **102** is configured for long-term storage of information as non-volatile memory space and retains information after power on/off cycles. In at least one implementation, the NVM **106** consists of one or more dies of NAND flash memory. Other examples of NVM **106** may include phase change memories, ReRAM memories, MRAM memories, magnetic media (including shingle magnetic recording), optical disks, floppy disks, electrically programmable read only memories (EPROM), electrically erasable programmable read only memories (EEPROM), and other solid-state memories. Magnetic media non-volatile memory

may be one or more magnetic platters in the data storage device **102**. Each platter may contain one or more regions of one or more tracks of data.

(12) The data storage device **102** includes a controller **110**, which manages operations of the data storage device **102**, such as writes to and reads from the NVM **106**. The controller **110** interfaces with the host **150** via mac and phy components **116**, and interfaces with the NVM **106** through a NVM interface **114**. The controller **110** includes a command parser **132** responsible for fetching and parsing commands from the host **150** and queuing them internally to a command accelerator **134**. The controller **110** includes one or more processors **130**, which may be multi-core processors. The processors **130** handle the components of the data storage device **102** through firmware code. The controller **110** further contains LBAs that are mapped or linked to physical block addresses (PBAs) of the storage elements in the NVM **106**.

(13) The controller **110** also includes volatile memory or cache **112** for short-term storage or temporary memory during operation of the data storage device **102**. The cache **112** does not retain stored data if powered off. Examples of volatile memories include random access memories (RAM), dynamic random access memories (DRAM), static random access memories (SRAM), and other forms of volatile memories. The cache **112** includes a random accumulated buffer **120** for supporting random data streams, a sequential accumulated buffer **124** for supporting sequential data streams, and an overlap accumulated buffer **128** for supporting overlap data streams. Each data stream is comprised of one or more write commands received from the host **150**.

(14) FIG. 2 illustrates a method **200** of a target device or a data storage device receiving data from a host and pushing the data to a cache, according to one embodiment. Method **200** will be described in reference to storage system **100** of FIG. 1, but other systems are possible. One or more operations of method **200** may be performed by the controller **110** of FIG. 1 executing machine-executable instructions in a non-transitory machine readable medium by a computer, hardware, a processor (e.g., a microprocessor), and/or machine.

(15) In operation **202**, one or more new commands are fetched or received by the controller **110** from the host **150** in a sequence. The one or more new commands are fetched or received depending on the host protocol. The controller **110** fetches the commands in NVMe protocol whereas the controller **110** receives the commands from queuing from the host **150** in UFS protocol or SATA protocol. For ease of description, the term “receiving” will encompass fetching, receiving, and arrival.

(16) In operation **204**, the commands are parsed and detected as write commands, read commands, or erase commands by the command parser **132**. The command parser **132** may further ensure there are no errors in the commands. If one or more commands are detected as write commands, method **200** proceeds to operation **206**.

(17) In operation **206**, the controller **110** determines whether there is a cache collision of the one or more write commands. A cache collision occurs when a new write command is received having an LBA range that overlaps with an LBA range of data currently stored in the cache **112** and not stored in NVM **106**. A cache collision may be detected by a collision table inputted with the write command sequence in order of receipt by the data storage device **102** and inputted with an LBA range of the data associated with the write command. If the controller **110** detects a cache collision, method **200** proceeds to operation **208**.

(18) In operation **208** one or more of the write commands are classified as overlap write commands. In other words, when the LBA range on a new write command overlaps with LBAs stored in cache **112**, the data associated with the new write command is classified as an overlap write command. If one or more write commands are classified as overlap write commands, method **200** proceeds to operation **212**.

(19) If in operation **206** the controller determines there is no cache collision, method **200** proceeds to operation **210**. In operation **210**, one or more of the write commands are classified as either random write commands for random data or sequential write commands for sequential data by the

command parser **132**. For example, when a write command has a large size of associated data (such as sized 32 KB or more), the write command may be classified as a sequential write command. For example, when a write command has small size of associated data, the write command may be classified as a random write command. The command parser **132** may perform a locality detection in order to determine whether the write commands should be classified as random write commands or sequential write commands. After classifying the write commands as random or sequential, method **200** proceeds to operation **212**.

(20) In operation **212**, the controller **110** determines if one or more of the write commands are classified as overlap write commands. For an overlap write command, the controller **110** pushes or writes the data associated with the overlap write command to the overlap accumulated buffer **128** in operation **214**. If the controller **110** determines that one or more of the write commands are not classified as overlap write commands, method **200** proceeds to operation **216**.

(21) In operation **216**, the controller **110** determines if one or more of the write commands are classified as random write commands. For a random write command, the controller **110** pushes or writes the data associated with the random write command to the random accumulated buffer **120** in operation **218**. If the controller **110** determines that one or more of the write commands are not classified as random write commands, method **200** proceeds to operation **220**, where the controller **110** pushes or writes the data associated with a sequential write command to the sequential accumulated buffer **124**. For each stream, the write data associated with the write commands is accumulated separately in the overlap, random, and sequential buffers.

(22) In operation **222**, when data associated with the one or more write commands are written to one of the separate buffers **120**, **124**, **128**, a command completion message is posted to the host **150**.

(23) Embodiments of method **200** permit a target device, such as the data storage device **102**, to receive any type of write commands without any performance degradation faced by the host **150**. The data storage device **102** is operable to keep receiving new data commands from the host **150** and to keep posting command completion to the host **150** even in the event of overlap data or write cache collision detection.

(24) FIG. **3** is a method **300** of writing data from a cache to NVM, according to one embodiment. Method **300** will be described in reference to system **100** of FIG. **1** with cache **112** including the random accumulated buffer **120**, sequential accumulated buffer **124**, and overlap accumulated buffer **128**, but other systems are possible. Method **300** may be used in conjunction with method **200** of FIG. **2**. One or more operations of method **300** may be performed by the controller **110** of FIG. **1** executing machine-executable instructions in a non-transitory machine readable medium by a computer, hardware, a processor (e.g., a microprocessor), and/or machine.

(25) In operation **302**, the controller **110** detects that one or more of the accumulated buffers **120**, **124**, **128** are available to be emptied to the NVM **106**. In at least one implementation, an accumulated buffer becomes available after accumulating enough data, such as by meeting or crossing a pre-configured threshold of data, when having a timeout or after a predetermined amount of time has passed, or when a flush request is received. In one embodiment, the threshold for making an accumulated buffer available is 512 KB or more.

(26) In operation **304**, the available accumulated buffer identified in operation **302** is queued in a hardware engine (such as for magnetic disk NVM) or to firmware (such as for NAND flash NVM or other solid state NVM) to be written to the NVM **106**.

(27) In operation **306**, the controller **110** determines if the available accumulated buffer queued in operation **304** is an accumulated buffer associated with an overlap stream, such as the overlap accumulated buffer **128**. If the controller **110** determines that the available accumulated buffer is not associated with an overlap stream, and is an accumulated buffer associated with a random stream or a sequential stream, such as the random accumulated buffer **120** or the sequential accumulated buffer **124**, the method **300** proceeds directly to operation **312**. In operation **312**, the

data in the available accumulated buffer is emptied to the NVM **106**. After writing the flushed data to the NVM **106**, the relevant LBAs are removed from a collision table, such as the collision table of operation **206** of FIG. 2.

(28) If in operation **306** the controller **110** determines that the available accumulated buffer is associated with an overlap stream, the method **300** proceeds to operation **308**. In operation **308**, the emptying of the available overlap accumulated buffer is temporarily delayed. The emptying of the available overlap accumulated buffer is delayed for so long as it takes the controller **110** to complete operation **310**.

(29) In operation **310**, the controller **110** flushes to the NVM **106** data in random streams and sequential streams that was received in a write command sequence prior to the data accumulated in the available overlap accumulated buffer. In other words, all random and sequential write commands that are received prior in the sequence to any overlap write command stored in the available overlap accumulated buffer are flushed to the NVM **106**. The random and sequential write commands are flushed to the NVM **106** even if the associated random accumulated buffer and sequential accumulated buffer are not currently available. The method **300** then proceeds to operation **312**, where the data in the available overlap accumulated buffer is emptied to the non-volatile memory. After writing the flushed data to the NVM **106**, the relevant LBAs are removed from a collision table, such as the collision table of operation **206** of FIG. 2.

(30) Thus, random write commands and sequential write commands received in prior sequence order to the overlap write commands stored in the available overlap accumulated buffer are first flushed to the NVM **106** before flushing the overlap write commands. As such, data from the overlap write commands stored in the available accumulated buffer will override the prior data commands. When the available overlap accumulated buffer is flushed to the NVM **106**, the overlap write commands are written to the NVM **106** in the order of receipt by target (i.e., in the order of a write command sequence) to help ensure that new write commands overrides prior write commands.

(31) Additional commands may continue to be received and accumulated in the random accumulated buffer, the sequential accumulated buffer, and the overlap accumulated buffer as method **300** is performed. However, only the data associated with random and sequential commands that were received prior in the write command sequence to the overlap commands in the available overlap accumulated buffer are flushed to the NVM **106**. The additional commands are queued later in sequence.

(32) From the point of view of the host **150**, the execution of overlap write commands is identical to the execution of random write commands or sequential write commands. This prevents the host **150** from facing any performance degradation from the data storage device **102** while executing the overlap commands.

(33) Embodiments of method **300** permit the data storage device **102** to handle overlap write data in compliance with host protocols to ensure that prior write commands do not override new write commands. Method **300** is operable to handle overlap write data without any performance degradation faced by the host **150**.

(34) Certain embodiments disclosed herein relate to a target device using a separate buffer to cache data associated with overlap data and to flush this separate data while ensuring that new write commands overrides prior write commands. Certain embodiments disclosed herein keep a pipeline open between a host and a target device without having performance degradation of the pipeline even in overlap write scenarios.

EXAMPLES

(35) The following are examples to illustrate various embodiments of handling write overlap data streams. For illustration purposes, the examples will be described in reference to FIGS. 1-3. Such examples are not meant to limit the scope of the disclosure unless specifically set forth in the claims.

Example 1

(36) In TABLE 1, a sequence of eleven write commands were queued in the cache **112** of the data storage device **102** in the order of receipt from the host **150**, and chronicled in the write command sequence 0 to 10. In this example, the data of each write command is associated with an LBA or LBA range, random data has a size of up to about 4 KB, sequential data has a size of 32 KB or more, and a threshold for making a buffer available is 512 KB or more, but other parameters are possible.

(37) Per method **200**, the first eight commands (write command sequence 0-7) received are classified as random write commands based on size and no overlap was detected in respect to previously queued write commands. As such, the write data of the first eight commands is accumulated in the random accumulated buffer **120**. The ninth command (write command sequence 8) is classified as an overlap write command since there is an LBA range overlap between this command and the seventh command as long as the data of the seventh command has not yet been flushed to the NVM **106**. The write data of the ninth command is accumulated in the overlap accumulated buffer **128**.

(38) The tenth command (write command sequence 9) is classified as a sequential write command because the size of the command is 32 KB and no overlap was detected in respect to previously queued write commands. The write data of the tenth command is accumulated in the sequential accumulated buffer **124**. The eleventh command (write command sequence 10) is classified as an overlap write command since there is an LBA range overlap between this command and the first command, the ninth command, and the tenth command, as long as the data of first, ninth, and tenth commands have not yet been flushed to the NVM **106**. The write data of the eleventh command is accumulated in the overlap accumulated buffer **128**.

(39) TABLE-US-00001

TABLE 1	Write command sequence	LBA range	Command classification
0	66	Random 1 100	Random 2 3000
1	Random 3 300	Random 4 406	Random 5 587
2	Random 6 6	Random 7 7870	Random 8 5-7
3	Overlap 9 128-135	Sequential 10 7-134	Overlap

(40) Per method **300**, the eleventh command (write command sequence 10) causes the overlap accumulated buffer to cross the threshold of 512 KB. Prior to flushing the overlap accumulated buffer to the NVM **106**, the data associated with the first write command through the eighth write command and the tenth write command (write command sequences 0-7 and 9) are flushed to the NVM **106** in the order received from the host **150** to empty the random accumulated buffer **120** and the sequential accumulated buffer **124**. The data associated with overlap accumulated buffer **128** is then flushed to the NVM **106**.

(41) Since the data associated with the seventh command (write command sequence 6) is flushed to the NVM **106** prior to the data associated with the ninth command (write command sequence 8), the new data of the ninth command overrides the prior data of the seventh command.

(42) Likewise, since the data associated with the first command (write command sequence 0), the ninth command (write sequence 8), and the tenth command (write command sequence 9) is flushed to the NVM **106** prior to the data associated with the eleventh command (write command sequence 10), the new data of the eleventh command overrides the prior data.

Example 2

(43) In a second example referring to TABLE 1, the random accumulated buffer **120** becomes available after receiving the sixth command (write command sequence 5). Per method **300**, once the controller **110** determines the available accumulated buffer is not an overlap accumulated buffer, the available random accumulated buffer may immediately be emptied to the NVM **106**. The controller **110** need not flush data associated with the overlap accumulated buffer or the sequential accumulated buffer prior to flushing the data associated with the random accumulated buffer.

(44) The host **150** does not face any performance degradation as the data storage device **102** executes all of the commands of TABLE 1. Additionally, the latest version of data received

overrides an earlier version of the data received previously.

(45) The methods and apparatus discussed above allow a target device or a data storage device to handle overlap write data in compliance with host protocols to ensure that prior write commands do not override new write commands. By utilizing a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer, cache coherency issues can be avoided without breaking the data pipeline. Thus, the storage system is able to handle overlap write data without any performance degradation faced by the host.

(46) In one embodiment, a data storage device comprises a non-volatile memory, and a controller coupled to the non-volatile memory. The controller comprises a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer. The controller is configured to receive a new write command, classify the new write command, and write data associated with the new write command to one of the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer based on the classification of the new write command.

(47) The data storage device may further comprise a command parser and/or a write command accelerator. The controller may be further configured to detect a cache collision of the new write command. The new write command may be classified as an overlap write command when a cache collision is detected. The random accumulated buffer may store one or more write commands classified as random, the sequential accumulated buffer may store one or more write commands classified as sequential, and the overlap accumulated buffer may store one or more write commands classified as overlap.

(48) In another embodiment, a storage system comprises a host device, and a data storage device coupled to the host device. The data storage device comprises a non-volatile memory, and a controller coupled to the non-volatile memory. The controller is configured to receive a new write command, push data associated to the new write command to one of a random stream, a sequential stream, or an overlap stream, and post a write command completion message to the host after the data associated with the new write command is pushed to one of the random stream, the sequential stream, or the overlap stream.

(49) The overlap stream may comprise one or more overlap accumulated buffers, the random stream may comprise one or more random accumulated buffers, and the sequential stream may comprise one or more sequential accumulated buffers. The controller may be further configured to empty the random stream, the sequential stream, or the overlap stream to the non-volatile memory after posting the write command completion message to the host. The controller may receive the new write command by fetching the new write command from the host. The controller may receive the new write command from the host queueing the new write command to the controller.

(50) In one embodiment, a method of handling write overlap data streams comprises receiving data to a controller in a sequence, detecting a cache collision of the data, and classifying the data as overlap data, random data, or sequential data. The method further comprises writing random data to a random buffer, writing sequential data to a sequential buffer, and writing overlap data to an overlap buffer, and emptying the overlap data buffer. Emptying the overlap data buffer comprises flushing the random data and the sequential data that were received prior to the overlap data in the sequence to a non-volatile memory, and flushing the overlap data to the non-volatile memory.

(51) Detecting the cache collision may comprise determining a logical block address range associated with the data overlaps with a logical block address range of data currently stored in the random buffer or the sequential buffer. Flushing the overlap data may overwrite a portion of the random data or a portion of the sequential data flushed to the non-volatile memory. Classifying the received data as overlap data may be based on the detected cache collision. Classifying the data as random data or sequential data may be based on a locality detection.

(52) In another embodiment, a method comprises detecting a first accumulated buffer of a controller is available to empty to a non-volatile memory, determining the first accumulated buffer is associated with an overlap stream, and delaying the emptying of the first accumulated buffer to

the non-volatile memory. The method further comprises flushing a second accumulated buffer associated with a random stream, flushing a third accumulated buffer associated with a sequential stream, and emptying the first accumulated buffer to the non-volatile memory.

(53) The first accumulated buffer may become available by meeting a pre-configured threshold of data. The first accumulated buffer may become available when a flush request is received. The first accumulated buffer may become available after a predetermined amount of time has passed. Write commands written to the second accumulated buffer or the third accumulated buffer may be received in a write command sequence prior to write commands written to the first accumulated buffer.

(54) The method may further comprise detecting a fourth accumulated buffer is available to empty to the non-volatile memory, determining the fourth accumulated buffer is associated with one of the random stream or the sequential stream, and flushing the fourth accumulated buffer to the non-volatile memory.

(55) In another embodiment, a data storage device comprises a non-volatile memory, means for detecting a cache collision, and means for ensuring that new write commands override prior write commands. The means for ensuring that new write commands override prior write commands comprises a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer.

(56) The non-volatile memory may comprise a NAND flash memory. The data storage device may further comprise means for classifying new write commands as random, sequential, or overlap. The data storage device may further comprise means for detecting when the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer are available to empty to the non-volatile memory.

(57) While the foregoing is directed to embodiments of the present disclosure, other and further embodiments of the disclosure may be devised without departing from the basic scope thereof, and the scope thereof is determined by the claims that follow.

Claims

1. A data storage device, comprising: a non-volatile memory; means for detecting a cache collision; means for ensuring that new write commands override prior write commands, wherein the means for ensuring that new write commands override prior write commands comprises a random accumulated buffer, a sequential accumulated buffer, and an overlap accumulated buffer; and a controller, wherein the controller is configured to: detect a first accumulated buffer of the controller is available to empty to the non-volatile memory; determine the first accumulated buffer is associated with an overlap stream; delay the emptying of the first accumulated buffer to the non-volatile memory; flush a second accumulated buffer associated with a random stream; flush a third accumulated buffer associated with a sequential stream; and empty the first accumulated buffer to the non-volatile memory.

2. The data storage device of claim 1, wherein the non-volatile memory comprises a NAND flash memory.

3. The data storage device of claim 1, further comprising means for classifying new write commands as random, sequential, or overlap.

4. The data storage device of claim 1, further comprising means for detecting when the random accumulated buffer, the sequential accumulated buffer, or the overlap accumulated buffer are available to empty to the non-volatile memory.

5. The data storage device of claim 1, wherein the controller comprises: the overlap accumulated buffer, the sequential accumulated buffer, and the random accumulated buffer.

6. The data storage device of claim 5, further comprising a write command accelerator.

7. The data storage device of claim 1, further comprising means to determine whether one or more

write commands are classified as overlap write commands.

8. The data storage device of claim 1, further comprising means to determine whether one or more write commands are classified as random write commands.

9. The data storage device of claim 1, further comprising means to write data associated with a sequential write commands to the sequential accumulated buffer.

10. The data storage device of claim 1, wherein the first accumulated buffer becomes available by meeting a pre-configured threshold of data.

11. The data storage device of claim 1, wherein the first accumulated buffer becomes available when a flush request is received.

12. The data storage device of claim 1, wherein the first accumulated buffer becomes available after a predetermined amount of time has passed.

13. The data storage device of claim 1, wherein write commands are received for writing to the second accumulated buffer or the third accumulated buffer, and wherein the write commands for writing to the second accumulated buffer or the third accumulated buffer are received in a write command sequence prior to write commands for writing to the first accumulated buffer.

14. The data storage device of claim 1, wherein the controller is further configured to: detect a fourth accumulated buffer is available to empty to the non-volatile memory; determine the fourth accumulated buffer is associated with one of the random stream or the sequential stream; and flush the fourth accumulated buffer to the non-volatile memory.

15. The data storage device of claim 1, wherein the controller is further configured to queue the accumulated buffers.

16. The data storage device of claim 1, wherein the emptying occurs after the flushing.

17. The data storage device of claim 1, wherein the delaying occurs prior to the flushing.

18. The data storage device of claim 1, wherein the delaying occurs after the determining.
