

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417181
Kind Code	B1
Date of Patent	September 16, 2025
Inventor(s)	Holey; Anup et al.

Systems and methods for aperture-specific cache operations

Abstract

A processing device including a first cache is coupled to a system memory and a parallel processing unit (PPU) including a second cache. An operation to modify cache lines of the second cache associated with a first aperture of the system memory is received. A first subset of cache lines of the second cache is identified. The first subset of cache lines is associated with the first aperture of the system memory and is different from a second subset of cache lines of a second aperture of the system memory. The first subset of cache lines is modified as specified by the cache operation.

Inventors: Holey; Anup (San Jose, CA), Gandhi; Wishwesh Anil (Sunnyvale, CA), Kaushikkar; Sujoyita (San Jose, CA), Mehra; Karan (Apex, NC), Robinson; Daniel Glenn (Cary, NC), Kiminki; Sami Olavi Johannes (Espoo, FI), Waterman; Alexander Michael (Westlake Hills, TX), Hairgrove; Mark (San Jose, CA), Smith; Jeff (Barrington, RI), Yin; Liang (Palo Alto, CA)

Applicant: NVIDIA Corporation (Santa Clara, CA)

Family ID: 1000007915801

Assignee: NVIDIA Corporation (Santa Clara, CA)

Appl. No.: 18/665392

Filed: May 15, 2024

Related U.S. Application Data

us-provisional-application US 63566142 20240315

Publication Classification

Int. Cl.: G06F12/0815 (20160101)

U.S. Cl.:

CPC **G06F12/0815** (20130101);

Field of Classification Search

USPC: None

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
11880327	12/2023	Nakibly	N/A	G06F 13/4027
2014/0049551	12/2013	Rao	345/543	G06T 1/20
2018/0293362	12/2017	Ray	N/A	H04N 19/00
2018/0293693	12/2017	Ray	N/A	G06F 12/0815
2020/0089611	12/2019	Gandhi	N/A	G06F 12/0808
2020/0192799	12/2019	Johns	N/A	G06F 12/1009

Primary Examiner: Patel; Kaushikkumar M

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

RELATED APPLICATIONS (1) This application claims the benefit of U.S. Provisional Patent Application No. 63/566,142 filed Mar. 15, 2024, the entire contents of which are incorporated by reference herein.

FIELD OF THE INVENTION

(1) Embodiments of the present disclosure generally relate to parallel processing systems. Specifically, embodiments of the present disclosure relate to parallel processing systems and methods for aperture-specific cache operations.

BACKGROUND

(2) Parallel processing in high-performance computing (HPC) systems involves the simultaneous execution of multiple computational tasks or operations. This is done by breaking down larger computations into smaller, independent subtasks that may be processed concurrently by multiprocessors. In some instances, parallel processing involves distributed computing such that task are distributed across multiple computing clusters. Each cluster may operate independently, and communication can be facilitated to share results.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) FIG. 1 is a block diagram illustrating a computer system configured to implement one or more aspects of the present disclosures.
- (2) FIG. 2 is a diagram illustrating an aperture-specific cache coherency system, in accordance with at least one embodiment of the present disclosure.

- (3) FIG. 3 illustrates a flowchart of a method for aperture-specific cache operations, in accordance with at least one embodiment of the present disclosure.
- (4) FIG. 4 illustrates a parallel processing unit, in accordance with at least one embodiment of the present disclosure.
- (5) FIG. 5A illustrates a general processing cluster within the parallel processing unit of FIG. 4, in accordance with at least one embodiment of the present disclosure.
- (6) FIG. 5B illustrates a memory partition unit of the parallel processing unit of FIG. 4, in accordance with at least one embodiment of the present disclosure.
- (7) FIG. 6A illustrates the streaming multi-processor of FIG. 5A, in accordance with at least one embodiment of the present disclosure.
- (8) FIG. 6B is a conceptual diagram of a processing system implemented using the PPU of FIG. 4, with at least one embodiment of the present disclosure.
- (9) FIG. 6C illustrates an exemplary system in which the various architecture and/or functionality of the various previous embodiments may be implemented.
- (10) FIG. 7 is a conceptual diagram of a graphics processing pipeline implemented by the PPU of FIG. 4, with at least one embodiment of the present disclosure.

DETAILED DESCRIPTION

- (11) High Performance Computing (HPC) systems may use specialized hardware architectures, such as parallel processing units (PPUs), to enhance parallel processing abilities. PPUs are designed to extract high performance using a large number of small, parallel execution threads on dedicated programmable multiprocessors. In PPUs, a group of threads, such as a warp, may execute the same instruction concurrently on a multiprocessor with different input data. This execution model is referred to as Single Instruction, Multiple Thread (SIMT) and is commonly utilized in parallel computing. PPUs are designed to execute a program (e.g., a kernel, a shader program, etc.) in parallel by executing many groups of threads on the PPU in which each thread of the groups of threads typically operates on a different portion of data. Because PPU architectures are generally optimized for executing many parallel threads simultaneously, PPUs, such as graphical processing units (GPUs), are leveraged to accelerate artificial intelligence (AI), HPC, cloud, and hyperscale workloads. Particularly, AI models are rapidly increasing in complexity and size as they enhance deep recommender systems containing large (e.g., 10 terabytes or more) amounts of data.
- (12) System architectures are emerging with fast access to memory enabled by a tight coupling of the central processing unit (CPU) and PPU via high-bandwidth interconnects, such as high-speed busses or specialized on-chip communication channels. Notably, bidirectional, high-bandwidth, and cache-coherent connections between CPU and PPU memory allow multiple application threads (CPU or PPU) to directly access system-allocated memory. For example, a PPU can cache system memory associated with and owned by the CPU within internal PPU caches. Similarly, the CPU can cache video memory associated with and owned by the PPU within internal CPU caches. Some conventional system may maintain coherency between GPU caches (e.g., L1, L2, L3, etc.) and CPU caches (e.g., L0, L1, L2, system level cache (SLC)) by software through explicit cache flushes and invalidates. For example, the CPU may allocate a portion of system memory and designate it as a buffer that the PPU can use for processing. When the PPU finishes its work on the buffer, there may be some data modified by the PPU associated with an address that belongs to the CPU's system memory but is known to the CPU. A software mechanism can ensure that the PPU has written out all of its data to the CPU cache before control of the buffer transfers back to the CPU. Cache invalidates may be carried out on the entire GPU cache such that every cache line associated with system memory is written back to the CPU cache and invalidated. This process can take many (e.g., thousands) of cycles and add latency between a GPU-CPU handshake. Such latency can hamper performance and reliability in latency sensitive systems such as Advanced Driver Assistance Systems (ADAS) and Automated Driving Systems (ADS).
- (13) To avoid issues associated with software-managed CPU-GPU coherency, some conventional

systems may utilize hardware coherency between GPU and CPU caches using a coherent memory interconnect with native hardware support. For example, chip-to-chip (C2C) hardware-coherency may enable a GPU to cache system memory within a cache at cache-line granularity without page migrations between system memory and video memory. Hardware coherence can simplify and speed up CPU-GPU communication for data closely shared between the CPU and the PPU. This hardware coherency can improve performance of memory accesses to non-local memory, such as a CPU thread accessing GPU memory (e.g., video memory) or a GPU thread accessing CPU memory (e.g., system memory). Hardware managed coherence simplifies the programming model as explicit software coherence is not used. For example, if a cache line is modified in a GPU cache, the CPU may receive an indication that the cache line has been modified and may retrieve associated data from the GPU when it is needed by the CPU. As such, explicit flushes and invalidates of GPU hardware coherent data may not be needed. However, coherently caching system memory buffers that may only be accessed by the PPU may incur needless overhead associated with hardware coherency as the CPU may never need access to data stored in such memory buffers.

(14) Aspects and implementations of the present disclosure address the above deficiencies and other deficiencies of conventional cache coherency systems by providing aperture-specific cache management operations for cache coherency. To enable memory sharing between a CPU and a PPU, the CPU may designate apertures associated with shared memory devices. An aperture is a portion of the address space that is persistently associated with a particular peripheral device or a memory unit. Apertures may reach external devices such as Read-Only Memory (ROM) or Random-Access Memory (RAM) chips, or internal memory on the CPU itself. For example, a memory device included in the system may begin addressing starting at zero. However, because the system has more than one memory device (e.g., system memory, video memory, etc.) with the same addressing scheme, the system may have ambiguous addressing. To resolve this, the system may designate multiple apertures, each associated with a memory device of the system. Thus, apertures form a layer of address translation below the level of virtual-physical mapping. For example, when a buffer is allocated in system memory for use by the PPU, the buffer may be designated as a system memory aperture. When a buffer is allocated in video memory, the buffer may be designated as a video memory aperture.

(15) In at least one embodiment, the system may maintain a logical distinction between non-coherent system memory aperture buffers and coherent system memory aperture buffers within the same system memory. Buffers designated as non-coherent system memory aperture buffers can include system memory buffers that the system expects only the PPU to access such as a game buffers, textures, compressible surfaces, and the like. Buffers designated as coherent system memory aperture buffers can include system memory buffers associated with workload sharing dispatched to the PPU that the system expects the CPU to access. When a system memory buffer is created, the system memory buffer can be designated as a non-coherent aperture buffer or a coherent aperture buffer. For example, if the buffer is regularly used to communicate between the CPU and the PPU, the buffer may be designated as a coherent system memory aperture buffer. If the buffer is expected to be used by the PPU and not the CPU, the buffer may be designated as a non-coherent system memory aperture buffer.

(16) In at least one embodiment, cache coherency can be managed based on aperture designations. Buffers designated as coherent system memory aperture buffers may be coherently cached according to hardware-managed coherence techniques. Coherency of buffers designated as non-coherent system memory aperture buffers may be managed using aperture-specific cache operations. In at least one embodiment, coherency of buffers designated as non-coherent memory aperture buffers may be managed through explicit cache flushes and invalidates. For example, before control of non-coherent memory buffer transfers from the PPU back to the CPU, a cache operation can be issued to flush and/or invalidate all cache lines designated as a non-coherent

system memory cache lines. The PPU can compare aperture identifiers in aperture fields of each cache line its cache and invalidate cache lines designated as non-coherent system memory aperture cache lines. Non-coherent system memory cache lines in one or more caches associated with the PPU can accordingly be flushed, written back to system memory, and invalidated. To invalidate cache lines, the PPU can writeback associated data back to CPU caches, mark the cache lines as invalid, and discard the cache lines. To flush cache lines, the PPU can writeback associated data back to CPU caches, and maintain the cache lines in a clean state for future reuse. It is appreciated that aperture-specific invalidate and flush cache operations are used herein by way of example, and not by way of limitation, noting that other cache operations can target aperture-specific cache lines.

(17) Advantages of the technology disclosure herein include, but are not limited, decreased destructive interference between coherent and non-coherent system memory cache lines. Specifically, by targeting specific aperture cache lines, latency associated with invalidating PPU cache lines can be reduced as fewer cache lines are targeted. Additionally, coherent system memory aperture cache lines that are critical to performance and are not associated with applications issuing non-coherent system memory aperture cache invalidates may be maintained within the cache and managed by hardware coherency techniques.

(18) FIG. 1 is a block diagram illustrating a computer system **100** configured to implement one or more aspects of the present disclosure. The system **100** includes a system memory **102** and a central processing unit (CPU) **104** that may communicate via an interconnection pathway such as a bus, a dedicated memory bridge, or other communication path. A parallel processing unit (PPU) **120** is operatively coupled to the CPU **104** via a communication path such as a system bus, a Peripheral Component Interconnect express (PCIe), a Northbridge/Southbridge architecture, or other communication path. In at least one embodiment, the PPU **120** may be integrated directly onto the CPU **104** die and communication may be handled internally within the CPU **104**. In yet another embodiment, the PPU **120** may be integrated with one or more other system elements, such as the CPU **104** and the system memory **102** to form a system on a chip (SoC). In such an embodiment, the CPU **104**, the PPU **120**, and other components of the system **100** may communicate using an architecture-specific common interface, such as an Advanced Reduced Instruction Set Computer (RISC) Machine (ARM) Common Hardware Interface (CHI) or a specialized Chip-to-Chip (C2C) interface.

(19) In at least one embodiment, the PPU **120** may incorporate circuitry optimized for graphics and video processing, including, for example, video output circuitry, and constitutes a graphics processing unit (GPU). In another embodiment, the PPU **120** incorporates circuitry optimizes for general purpose processing, while preserving the underlying computational architecture. The system **100** further includes a video memory (VMEM) **130** that the PPU **120** may use to store necessary data such as textures, frame buffers, shaders, and other graphical elements. In at least one embodiment, the VMEM **130** can include various types of memory devices, included dynamic random-access memory (DRAM) or graphics random-access memory such as video random-access memory (VRAM) or synchronous graphics random-access memory (SGRAM), including graphics double data rate (GDDR) SGRAM. In at least one embodiment, the VMEM **130** can include one or more stacks of memory, such as multiple DRAM dies stacked vertically, to form a high bandwidth memory (HBM). It is appreciated that the specific implementation of VMEM **130** can vary and can be selected from one of many available designs.

(20) The PPU **120** may include one or more multiprocessor(s) **124A** through **124N** (referred to generally as a multiprocessor(s) **124**, herein) and a cache **122**. Each multiprocessor **124**, for example, may be a streaming multiprocessor (SM), a compute unit (CU), a many integrated core (MIC), and the like. Each multiprocessor **124** may be responsible for executing parallel processing tasks, which involve performing the same operation on multiple pieces of data concurrently. Each multiprocessor **124** can execute a certain number of threads simultaneously such that the PPU **120**, as a whole, can execute a large number of threads concurrently across all multiprocessors **124**.

Each multiprocessor **124** may include an L1 cache (not shown) or use a corresponding L1 cache outside of the multiprocessor **124** that is used to perform load and store operations. Each multiprocessor **124** may have access to the cache **122** that is shared among all multiprocessors **124** and may be used to transfer data between threads. The cache **122** may be arranged along any level of a cache hierarchy (e.g., L1, L2, L3, etc.). In at least one embodiment, the cache **122** can be a shared memory (e.g., a local memory) that can behave as a programmable cache shared among threads. Each multiprocessor **124** may also have access to global memory, which can include, for example, system memory **102** and/or VMEM **130**.

(21) The CPU **104** can include one or more processing cores **106A** through **106N** (referred to generally as “processing cores **106**” herein) and a cache **110**. Each of the processing cores **106** may be individual processing units within the CPU **104** that independently execute instructions. Each processing core **106** may be a complete processing unit with its own arithmetic logic units (ALUs), control units, registers, and other components necessary to execute program instructions. Each processing core **106** may have dedicated cache (e.g., L1, L2, etc.) and may further have access to the cache **110**. In at least one embodiment, the cache **110** may be a shared cache such as an L3 cache, an L4 cache, a system level cache (SLC), etc. Each processing core **160** can additionally have access to global memory, such as system memory **102**.

(22) It is appreciated that the system **100** illustrated herein is illustrative and that variations and modifications are possible. The connection topology, the number of CPUs **104**, the number of PPU **120**, the number of processing cores **106** within the CPU **104**, and the number of multiprocessors **124** within the PPU **120** may be modified as desired. Additionally, the particular components illustrated herein are not exhaustive; for example, any number of add-in cards, peripheral devices, switches, network adapters, and the like might be supported but are not illustrated herein.

(23) In at least one embodiment, the PPU **120** is a graphics processor with rendering pipelines that can be configured to perform various tasks related to generated pixel data from graphics data supplied by CPU **104** and/or system memory. In operation, the CPU **104** is the central processor of computer system **100**, controlling and coordinating operations of other system components. In particular, the CPU **104** may issue commands that control the operation of PPU **120**.

(24) In at least one embodiment, the CPU **104** and the PPU **120** may be tightly coupled such that CPU **104** and the PPU **120** can share system memory **102** and VMEM **130**. The system **100** can include high-bandwidth interconnects, such as high-speed busses or on-chip communication channels, e.g., ARM CHI. Such a design is aimed at achieving efficient collaboration between the CPU **104** and PPU **120** for parallel processing tasks, such as graphics rendering or general-purpose computing.

(25) In some instances, the CPU may offload specific computational tasks from the CPU to the PPU, for example, to take advantage of parallel processing capabilities of the GPU such as graphics rendering, simulations, machine learning, and the like. Responsive to identifying specific tasks within a program suitable for execution on the PPU, the CPU may share relevant data stored in system memory **102** with the PPU **120**. For example, CPU **104** can set aside a memory block, such as a buffer, in system memory **102** for PPU **120** to use.

(26) In at least one embodiment, address translation service (ATS) can allow the CPU **104** and the PPU **120** to share one or more per-process page tables which enable both CPU **104** threads and PPU **120** threads to access system-allocated memory residing in physical memory such as system memory **102** or VMEM **130**. For example, chip-to-chip (C2C) hardware-coherency may enable PPU **120** to cache system memory **102** buffers within the cache **122** at cache-line granularity without page migrations between system memory **102** and VMEM **130**. Hardware coherence can simplify and speed up CPU-GPU communication for data (referred to as “coherent data” or “coherent buffers” herein) closely shared between the CPU **104** and the PPU **120**. In at least one embodiment, the system **100** may include a hardware interface protocol using a directory-based

system for hardware cache coherency between CPU **104** caches and PPU **120** caches. For example, the directory-based protocol may include several states for each cache line, such as Modified (M), Exclusive (E), Shared (S), and Invalid (I). However, coherently caching certain system memory buffers (referred to as “non-coherent data” or “non-coherent buffers” herein) that are only be accessed by the PPU **120** may incur needless overhead associated with hardware coherency. Accordingly, an aperture manager **112** of the CPU **104** and/or an aperture manager **121** of the PPU **120** may issue aperture-specific cache operations to manage sharing of non-coherent system memory buffers.

(27) The aperture managers **112** and **121** may be software component of a software sequence for managing communication between the CPU **104** and PPU **120**. The software sequence can include an application programming interface (API), device initialization, memory allocation, data transfer between CPU **104** and PPU **120**, kernel compilation and execution, CPU-GPU coherency, and the like. Specifically, the aperture manager **112** may enable sharing of non-coherent system memory buffers between the CPU **104** and PPU **120**. To enable memory sharing, the aperture manager **112** may designate apertures associated with shared memory devices. An aperture is a portion of the address space that is persistently associated with a particular peripheral device or a memory unit. Apertures may reach external devices such as Read-Only Memory (ROM) or Random-Access Memory (RAM) chips, or internal memory on the CPU **104** itself. For example, a memory device included in the system **100** may address starting at zero. However, because the system **100** has more than one memory device (e.g., system memory **102**, VMEM **130**, etc.), the system **100** would have ambiguous addressing. To resolve this, aperture manager **112** may designate one of multiple apertures when a buffer is allocated from a memory device. For example, when a buffer is allocated in system memory **102**, the aperture manager **112** may designate the buffer as a system memory aperture. When a buffer is allocated in VMEM **130**, the aperture manager **112** may designate the buffer as a video memory aperture.

(28) In at least one embodiment, the aperture manager **112** may maintain a distinction between a non-coherent system memory aperture and a coherent system memory aperture. Buffers designated as non-coherent system memory apertures can include system memory buffers that only the PPU **120** may access such as a game buffers, textures, compressible surfaces, and the like. Buffers designated as coherent system memory aperture buffers can include system memory buffers associated with workload sharing dispatch to the PPU **120**. When a system memory buffer is created, the aperture manager **112** can designate the buffer as a non-coherent system memory aperture or a coherent system memory aperture. For example, the aperture manager **112** may designate buffers regularly used to communicate between the CPU **104** and the PPU **120** as coherent system memory aperture buffers. The aperture manager **112** may designate buffers including data used by the PPU **120**, and not the CPU **104**, as a non-coherent system memory aperture buffers.

(29) In at least one embodiment, cache coherency can be managed based on aperture designations. Buffers designated as coherent system memory apertures may be coherently cached in the cache **122** according to hardware-managed coherency techniques described above. Coherency of buffers designated as non-coherent system memory aperture buffers may be managed using aperture-specific cache operations. In at least one embodiment, coherency of buffers designated as non-coherent memory aperture buffers may be maintained by the aperture manager **121** of the PPU **120** through explicit cache flushes and invalidates. For example, before control of non-coherent memory buffer transfers from the PPU **120** back to the CPU **104**, the aperture manager **121** may issue a cache-management operations to the cache **122** to flush and/or invalidate all cache lines associated with a non-coherent system memory aperture. Responsive to receiving such a cache operation, The PPU **120** can compare fields in each cache line of the cache **122** corresponding to aperture identifiers and selectively flush and/or invalidate cache lines associated with a non-coherent system memory aperture. Non-coherent system memory cache lines in the cache **122** can

accordingly be flushed from the cache, written back to system memory **102** and/or one or more CPU **104** caches, and invalidated.

(30) FIG. **2** is a diagram **200** illustrating an aperture-specific cache coherency system, in accordance with at least one embodiment of the present disclosure. Diagram **200** may include similar elements illustrated by computing system **100**, as described with respect to FIG. **1**. It can be noted that elements of FIG. **1** can be used herein to help describe FIG. **2**. The operations described with respect to FIG. **2** are described to be performed serially for the sake of illustration, rather than limitation. Although shown in a particular sequence or order, unless otherwise specified, the order of operations can be modified. Thus, the illustrated embodiments should be understood only as examples, and the described operations can be performed in a different order, while some operations can be performed in parallel. Additionally, one or more operations may be omitted in some embodiments. Thus, not all described operations are required in every embodiment, and other process flows are possible. In some embodiments, the same, different, fewer, or greater operations may be performed. Diagram **200** illustrates a technique that utilizes memory aperture buffer designations as part of a cache coherency technique. The diagram **200** illustrates a CPU **104** including a cache **122** and an aperture manager **112**, a system memory associated with the CPU **104**, a PPU **120** including a cache **122**, an aperture manager **121**, and a VMEM **130** associated with the PPU **120**.

(31) The cache **122** may be configured to cache data stored at system memory **102** and/or VMEM **130**. In at least one embodiment, the CPU **104** may allocate a buffer in system memory **102** to be processed by the aperture manager **112**. The size of the allocated system memory buffer can depend on the amount of data to be processed by the PPU **120**. The CPU **104** may populate the system memory buffer with necessary data and transfer the buffer from the CPU **104** to the PPU **120**. In at least one embodiment, prior to transferring the system memory buffer to the PPU **120**, the aperture manager **112** can designate an aperture associated with the system memory. In at least one embodiment, the aperture manager **112** may designate the system memory buffer as a non-coherent system memory buffer or a coherent system memory buffer. In at least one embodiment, the aperture manager **112** may designate buffers containing data that are primarily intended to be accessed and used by the PPU **120** as non-coherent system memory aperture buffers. For example, the aperture manager **112** may designate buffers containing texture data, vertex data, shader programs, compressible surfaces, framebuffer data, gaming buffers, GPU buffers, and the like as non-coherent system memory aperture buffers. In at least one embodiment, the aperture manager **112** can designate buffers (e.g., vertex buffers) containing data that are intended to be closely shared between the PPU **120** and the CPU **104** as coherent system memory aperture buffers. For example, if the system memory buffer is regularly used to communicate between the CPU and the PPU, the aperture manager **112** may designate the buffer as a coherent system memory aperture buffer.

(32) When the PPU **120** writes data to the cache **122**, the cache **122** can be updated based on apertures associated with the data, as designated by the aperture manager **112**. In an illustrative example, the cache **122** can include multiple cache lines **208A** through **208F** (referred to generally as “cache lines **208**” herein) where each cache line **208** includes at least an aperture field **202** and a data field **206**. The aperture field **202** may include one or more bits indicating an aperture associated with the data **206**. For example, ‘01’ may indicate that the cached data is associated with a non-coherent system memory aperture buffer, ‘10’ may indicate that the cached data is associated with a coherent system memory aperture buffer, and ‘11’ may indicate that the cached data is associated with a video memory aperture buffer. As such, data stored at cache line **208A**, cache line **208C**, and cache line **208D** may be associated with a non-coherent system memory aperture buffer; data stored at cache line **208B** may be associated with a coherent system memory aperture buffer; and data stored at cache line **208F** may be associated with a video memory aperture buffer.

(33) The aperture manager **121** may issue aperture-specific cache line invalidate operations to the

cache **122** after PPU **120** has finished working on the allocated buffer. For example, the aperture manager **121** can issue a cache operation to invalidate cache lines in the cache **122** associated with a non-coherent system memory aperture. Responsive to reception of a cache operation to invalidate non-coherent system memory aperture cache lines, a cache controller (not illustrated) of the cache **122** can sequentially access each cache line **208** to determine identifiers stored at respective aperture fields **202** of the cache lines **208**. The cache controller can further cause cache lines storing data associated with a non-coherent system memory aperture to be flushed and invalidated. For example, responsive to receiving a cache operation to invalidate cache lines associated with a non-coherent memory aperture, cache lines **208A**, **208C**, and **208D** can be flushed and invalidated. (34) To flush cache lines associated with a non-coherent memory aperture, the data may be written back to system **102** and/or one or more caches associated with the CPU **104**, such as cache **122**. In at least one embodiment, the cache **122** is a write-back cache such that write back of modified data to system memory **102** occurs when the cache line is flushed and invalidated as a result of a cache operation to invalidate non-coherent system memory cache lines. In at least one embodiment, to invalidate non-coherent system memory aperture cache lines, a state of the cache lines can be updated to reflect invalid status. In at least one embodiment, the cache **122** may operate according to a Modified, Exclusive, Shared, Invalid (MESI) coherency protocol such that one or bits of the non-coherent memory aperture cache lines are updated to reflect an invalid state. For example, a valid bit associated with the cache lines **208A**, **208C**, and **208D** can be updated to reflect an invalid state.

(35) In at least one embodiment, the PPU **120** can ensure that coherent system memory caches lines remain in the cache **122** and are not invalidated by operations to invalidate non-coherent system memory aperture cache lines. For example, in response to aperture manager **121** issuing a cache operation to invalidate non-coherent system memory cache lines, a cache controller of the cache **122** can ensure that cache line **208B** remains in the cache **122** and is not invalidated. By using software-managed non-coherent system memory aperture cache line invalidates, coherent system memory aperture cache lines can be managed using hardware coherency protocols without destructive interference between coherent and non-coherent system memory cache lines in the cache **122**.

(36) In at least one embodiment, cache lines **208** can include one or more fields not illustrated with respect to FIG. 2. For example, cache lines **208** can each include a tag, an index, an offset, a block offset, etc. In at least one embodiment, cache lines **208** can include one or more additional fields to enable software-managed coherency at varying levels of abstraction. For example, cache lines **208** can include an additional field indicating a process associated with the data. The CPU **104** can be a multicore system with multiple processes that operate independently. A first process executing on the CPU **104** may allocate a first non-coherent buffer of data to be processed by the PPU **120**. A second process executing on the CPU **104** may allocate a second non-coherent buffer of data to be processed by the PPU **120**. The Cache **122** can write a portion of the first non-coherent buffer of data to cache line **208A** and update a field of the cache line **208A** with an identifier of the first process. The cache **122** may write a portion of the second non-coherent buffer of data to cache lines **208C** and **208D** and update respective fields with an identifier of the second process. When the CPU **104** is prepared to resume control of the first non-coherent buffer, the aperture manager **112** may initiate a cache operation to invalidate all non-coherent cache lines associated with the first process. Resultantly, cache line **208A** may be flushed and invalidated while cache lines **208C** and **208D** remain valid and in the cache **122**.

(37) In at least one embodiment, the aperture manager **121** may issue non-coherent aperture cache line invalidates upon completion of associated tasks accelerated by the PPU **120**. For example, an application processing on the CPU **104** may request (e.g., using aperture manager, an API, etc.) a non-coherent system memory **102** buffer for GPU. An operating system (OS) running on the CPU **104** may allocate the non-coherent system memory buffer and notify the PPU **120**. For example, a

device driver associated with the PPU **120** may interface between the OS and PPU **120** hardware such that the device driver is notified when the OS allocates memory PPU **120** usage. Once the non-coherent system memory buffer is allocated and transferred to the PPU **120**, the PPU **120** can launch kernels or tasks to perform computations on the PPU **120** using its parallel processing capabilities. After executing the assigned kernels/tasks, the PPU **120** may notify (e.g., using interrupts, events, interprocess communication, and/or the like) the CPU **104** that the assigned tasks have been completed. Upon receiving an indication that the PPU **120** has completed the assigned tasks, the PPU **120** may issue a cache line operation to invalidate non-coherent system memory aperture cache lines.

(38) FIG. **3** illustrates a flowchart of a method **300** for aperture-specific cache operations, in accordance with at least one embodiment of the present disclosure. Although method **300** is described in the context of a processing unit, the method **300** may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method **300** may be executed by a GPU (graphics processing unit), CPU (central processing unit), or any processor capable of issuing or receiving aperture-specific cache operations. Furthermore, persons of ordinary skill in the art will understand that any system that performs method **300** is within the scope and spirit of embodiments of the present invention.

(39) More illustrative information will now be set forth regarding various optional architectures and features with which the foregoing framework may be implemented, per the desires of the user. It should be strongly noted that the following information is set forth for illustrative purposes and should not be construed as limiting in any manner. Any of the following features may be optionally incorporated with or without the exclusion of other features described.

(40) Although shown in a particular sequence or order, unless otherwise specified, the order of the operations can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated operations can be performed in a different order, and some operations can be performed in parallel. Additionally, one or more operations can be omitted in various embodiments. Thus, not all operations are required in every embodiment.

(41) At operation **302** of method **300**, processing circuitry may receive, at a parallel processing unit (PPU) including a first cache, a cache operation to modify cache lines of the first cache associated with a first aperture of a system memory of a processing device. In at least one embodiment, the processing device can be CPU **104** of FIG. **1**, the first cache can be cache **122** of FIG. **1**, and the PPU can be PPU **120** of FIG. **1**. In at least one embodiment, the PPU and the processing device are interconnected via an interface using a common hardware interface (CHI) protocol.

(42) At operation **304** of method **300**, the processing circuitry may identify a first subset of cache lines of the first cache, where the first subset of cache lines is associated with the first aperture of the system memory. In an illustrative example, the first subset of cache lines may include cache lines **208A**, **208C**, and **208D** of FIG. **2**. In at least one embodiment, the first aperture is a non-coherent aperture of the system memory.

(43) At operation **306** of method **300**, the processing circuitry may identify a second subset of cache lines of the first cache, where the second subset of cache lines is associated with a second aperture of the system memory. In an illustrative example, the second subset of cache lines may include cache line **208B** of FIG. **2**. In at least one embodiment, the second aperture is a coherent aperture of the system memory. In at least one embodiment, the first subset of cache lines and the second subset of cache lines are identified based on identifiers associated with the non-coherent aperture and the coherent aperture, respectively.

(44) At operation **308** of method **300**, the processing circuitry may modify the first subset of cache lines as specified by the cache operation. For example, the processing circuitry may modify cache lines **208A**, **208C**, and **208D** of FIG. **2**. In at least one embodiment, the cache operation is an invalidate operation, and to modify the first subset of cache lines the processing circuitry is to write data stored at the first subset of cache lines back to a second cache of the processing device, and

invalidate the first subset of cache lines. In at least one embodiment, the second cache may be cache **110** of FIG. **1**. In at least one embodiment, the cache operation is a flush operation, and to modify the first subset of cache lines, the processing circuitry is to write data stored at the first subset of cache lines back to the second cache of the processing device, and maintain a clean state of the first subset of cache lines.

(45) In at least one embodiment, the processing circuitry may further cause the second subset of cache lines to be maintained within the second cache. In at least one embodiment, coherency of the second subset of cache lines is managed by hardware associated with the processing device. For example, coherency cache line associated with a coherent system memory aperture may be managed based on a hardware interface using a directory-based approach.

(46) In at least one embodiment, the first subset of cache lines can further be differentiated and invalidated based on a process identifier that indicates one of a plurality of processes associated with the processing device. For example, the processing circuitry may receive an operation to invalidate non-coherent aperture cache lines of the second cache associated with a first process. The process logic may further invalidate one or more cache lines of the first subset of cache lines associated with the first process.

Parallel Processing Architecture

(47) FIG. **4** illustrates a parallel processing unit (PPU) **400**, in accordance with an embodiment. In an embodiment, the PPU **400** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The PPU **400** is a latency hiding architecture designed to process many threads in parallel. A thread (e.g., a thread of execution) is an instantiation of a set of instructions configured to be executed by the PPU **400**. In an embodiment, the PPU **400** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device such as a liquid crystal display (LCD) device. In other embodiments, the PPU **400** may be utilized for performing general-purpose computations. While one exemplary parallel processor is provided herein for illustrative purposes, it should be strongly noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

(48) One or more PPUs **400** may be configured to accelerate thousands of High Performance Computing (HPC), data center, and machine learning applications. The PPU **400** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image, and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and the like.

(49) As shown in FIG. **4**, the PPU **400** includes an Input/Output (I/O) unit **405**, a front-end unit **415**, a scheduler unit **420**, a work distribution unit **425**, a hub **430**, a crossbar (Xbar) **470**, one or more processing clusters **450** (e.g., general processing clusters (GPCs), and one or more partition units **480**. The PPU **400** may be connected to a host processor or other PPUs **400** via one or more high-speed NVLink **410** interconnect. The PPU **400** may be connected to a host processor or other peripheral devices via an interconnect **402**. The PPU **400** may also be connected to a local memory comprising a number of memory devices **404**. In an embodiment, the local memory may comprise a number of dynamic random-access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device.

(50) The NVLink **410** interconnect enables systems to scale and include one or more PPUs **400** combined with one or more CPUs, supports cache coherence between the PPUs **400** and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **410** through the hub

430 to/from other units of the PPU **400** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). The NVLink **410** is described in more detail in conjunction with FIG. **4B**.

(51) The I/O unit **405** is configured to transmit and receive communications (e.g., commands, data, etc.) from a host processor (not shown) over the interconnect **402**. The I/O unit **405** may communicate with the host processor directly via the interconnect **402** or through one or more intermediate devices such as a memory bridge. In an embodiment, the I/O unit **405** may communicate with one or more other processors, such as one or more the PPUs **400** via the interconnect **402**. In an embodiment, the I/O unit **405** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect **402** is a PCIe bus. In alternative embodiments, the I/O unit **405** may implement other types of well-known interfaces for communicating with external devices.

(52) The I/O unit **405** decodes packets received via the interconnect **402**. In an embodiment, the packets represent commands configured to cause the PPU **400** to perform various operations. The I/O unit **405** transmits the decoded commands to various other units of the PPU **400** as the commands may specify. For example, some commands may be transmitted to the front-end unit **415**. Other commands may be transmitted to the hub **430** or other units of the PPU **400** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **405** is configured to route communications between and among the various logical units of the PPU **400**.

(53) In an embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the PPU **400** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer is a region in a memory that is accessible (e.g., read/write) by both the host processor and the PPU **400**. For example, the I/O unit **405** may be configured to access the buffer in a system memory connected to the interconnect **402** via memory requests transmitted over the interconnect **402**. In an embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the PPU **400**. The front-end unit **415** receives pointers to one or more command streams. The front-end unit **415** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the PPU **400**.

(54) The front-end unit **415** is coupled to a scheduler unit **420** that configures the various processing clusters **450** to process tasks defined by the one or more streams. The scheduler unit **420** is configured to track state information related to the various tasks managed by the scheduler unit **420**. The state may indicate which processing cluster **450** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **420** manages the execution of a plurality of tasks on the one or more processing clusters **450**.

(55) The scheduler unit **420** is coupled to a work distribution unit **425** that is configured to dispatch tasks for execution on the processing clusters **450**. The work distribution unit **425** may track a number of scheduled tasks received from the scheduler unit **420**. In an embodiment, the work distribution unit **425** manages a pending task pool and an active task pool for each of the processing clusters **450**. The pending task pool may comprise a number of slots (e.g., 32 slots) that contain tasks assigned to be processed by a particular processing cluster **450**. The active task pool may comprise a number of slots (e.g., 4 slots) for tasks that are actively being processed by the processing clusters **450**. As a processing cluster **450** finishes the execution of a task, that task is evicted from the active task pool for the processing cluster **450** and one of the other tasks from the pending task pool is selected and scheduled for execution on the processing cluster **450**. If an active task has been idle on the processing cluster **450**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the processing cluster **450** and returned to the pending task pool while another task in the pending task pool is selected and scheduled for execution on the processing cluster **450**.

(56) The work distribution unit **425** communicates with the one or more processing clusters **450** via XBar **470**. The XBar **470** is an interconnect network that couples many of the units of the PPU **400** to other units of the PPU **400**. For example, the XBar **470** may be configured to couple the work distribution unit **425** to a particular processing cluster **450**. Although not shown explicitly, one or more other units of the PPU **400** may also be connected to the XBar **470** via the hub **430**.

(57) The tasks are managed by the scheduler unit **420** and dispatched to a processing cluster **450** by the work distribution unit **425**. The processing cluster **450** is configured to process the task and generate results. The results may be consumed by other tasks within the processing cluster **450**, routed to a different processing cluster **450** via the XBar **470**, or stored in the memory **404**. The results can be written to the memory **404** via the partition units **480**, which implement a memory interface for reading and writing data to/from the memory **404**. The results can be transmitted to another PPU **400** or CPU via the NVLink **410**. In an embodiment, the PPU **400** includes a number *U* of partition units **480** that is equal to the number of separate and distinct memory devices **404** coupled to the PPU **400**. A partition unit **480** will be described in more detail below in conjunction with FIG. 5B.

(58) In an embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the PPU **400**. In an embodiment, multiple compute applications are simultaneously executed by the PPU **400** and the PPU **400** provides isolation, quality of service (QoS), and independent address spaces for the multiple compute applications. An application may generate instructions (e.g., API calls) that cause the driver kernel to generate one or more tasks for execution by the PPU **400**. The driver kernel outputs tasks to one or more streams being processed by the PPU **400**. Each task may comprise one or more groups of related threads, referred to herein as a warp. In an embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including instructions to perform the task and that may exchange data through shared memory. Threads and cooperating threads are described in more detail in conjunction with FIG. 6A.

(59) FIG. 5A illustrates a processing cluster **450** of the PPU **400** of FIG. 4, in accordance with an embodiment. As shown in FIG. 5A, each processing cluster **450** includes a number of hardware units for processing tasks. In an embodiment, each processing cluster **450** includes a pipeline manager **510**, a pre-raster operations unit (PROP) **515**, a raster engine **525**, a work distribution crossbar (WDX) **580**, a memory management unit (MMU) **590**, and one or more Data Processing Clusters (DPCs) **520**. It will be appreciated that the processing cluster **450** of FIG. 5A may include other hardware units in lieu of or in addition to the units shown in FIG. 5A.

(60) In an embodiment, the operation of the processing cluster **450** is controlled by the pipeline manager **510**. The pipeline manager **510** manages the configuration of the one or more DPCs **520** for processing tasks allocated to the processing cluster **450**. In an embodiment, the pipeline manager **510** may configure at least one of the one or more DPCs **520** to implement at least a portion of a graphics rendering pipeline. For example, a DPC **520** may be configured to execute a vertex shader program on the programmable multiprocessor **540**. Multiprocessors **540** may generally include streaming multiprocessors, compute units, many integrated cores, and the like. The pipeline manager **510** may also be configured to route packets received from the work distribution unit **425** to the appropriate logical units within the processing cluster **450**. For example, some packets may be routed to fixed function hardware units in the PROP **515** and/or raster engine **525** while other packets may be routed to the DPCs **520** for processing by the primitive engine **535** or the multiprocessor **540**. In an embodiment, the pipeline manager **510** may configure at least one of the one or more DPCs **520** to implement a neural network model and/or a computing pipeline.

(61) The PROP unit **515** is configured to route data generated by the raster engine **525** and the DPCs **520** to a Raster Operations (ROP) unit, described in more detail in conjunction with FIG. 5B. The PROP unit **515** may also be configured to perform optimizations for color blending, organize

pixel data, perform address translations, and the like.

(62) The raster engine **525** includes a number of fixed function hardware units configured to perform various raster operations. In an embodiment, the raster engine **525** includes a setup engine, a coarse raster engine, a culling engine, a clipping engine, a fine raster engine, and a tile coalescing engine. The setup engine receives transformed vertices and generates plane equations associated with the geometric primitive defined by the vertices. The plane equations are transmitted to the coarse raster engine to generate coverage information (e.g., an x,y coverage mask for a tile) for the primitive. The output of the coarse raster engine is transmitted to the culling engine where fragments associated with the primitive that fail a z-test are culled, and transmitted to a clipping engine where fragments lying outside a viewing frustum are clipped. Those fragments that survive clipping and culling may be passed to the fine raster engine to generate attributes for the pixel fragments based on the plane equations generated by the setup engine. The output of the raster engine **525** comprises fragments to be processed, for example, by a fragment shader implemented within a DPC **520**.

(63) Each DPC **520** included in the processing cluster **450** includes an M-Pipe Controller (MPC) **530**, a primitive engine **535**, and one or more Multiprocessors **540**. The MPC **530** controls the operation of the DPC **520**, routing packets received from the pipeline manager **510** to the appropriate units in the DPC **520**. For example, packets associated with a vertex may be routed to the primitive engine **535**, which is configured to fetch vertex attributes associated with the vertex from the memory **404**. In contrast, packets associated with a shader program may be transmitted to the multiprocessor **540**.

(64) In some embodiments, the multiprocessor **540** comprises a programmable multiprocessor, such as a programmable streaming multiprocessor, that is configured to process tasks represented by a number of threads. Each multiprocessor **540** is multi-threaded and configured to execute a plurality of threads (e.g., 32 threads) from a particular group of threads concurrently. In an embodiment, the multiprocessor **540** implements a SIMD (Single-Instruction, Multiple-Data) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on the same set of instructions. All threads in the group of threads execute the same instructions. In another embodiment, the multiprocessor **540** implements a SIMT (Single-Instruction, Multiple Thread) architecture where each thread in a group of threads is configured to process a different set of data based on the same set of instructions, but where individual threads in the group of threads are allowed to diverge during execution. In an embodiment, a program counter, call stack, and execution state are maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call stack, and execution state are maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency. The multiprocessor **540** will be described in more detail below in conjunction with FIG. **6A**.

(65) The MMU **590** provides an interface between the processing cluster **450** and the partition unit **480**. The MMU **590** may provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In an embodiment, the MMU **590** provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory **404**.

(66) FIG. **5B** illustrates a memory partition unit **480** of the PPU **400** of FIG. **4**, in accordance with an embodiment. As shown in FIG. **5B**, the memory partition unit **480** includes a Raster Operations (ROP) unit **550**, a level two (L2) cache **560**, and a memory interface **570**. The memory interface **570** is coupled to the memory **404**. Memory interface **570** may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. In an embodiment, the PPU **400** incorporates U memory interfaces **570**, one memory interface **570** per pair of partition units **480**, where each pair

of partition units **480** is connected to a corresponding memory device **404**. For example, PPU **400** may be connected to up to Y memory devices **404**, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random-access memory, or other types of persistent storage.

(67) In an embodiment, the memory interface **570** implements an HBM2 memory interface and Y equals half U. In an embodiment, the HBM2 memory stacks are located on the same physical package as the PPU **400**, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In an embodiment, each HBM2 stack includes four memory dies and Y equals 4, with HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

(68) In an embodiment, the memory **404** supports Single-Error Correcting Double-Error Detecting (SECCDED) Error Correction Code (ECC) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where PPUs **400** process very large datasets and/or run applications for extended periods.

(69) In an embodiment, the PPU **400** implements a multi-level memory hierarchy. In an embodiment, the memory partition unit **480** supports a unified memory to provide a single unified virtual address space for CPU and PPU **400** memory, enabling data sharing between virtual memory systems. In an embodiment the frequency of accesses by a PPU **400** to memory located on other processors is traced to ensure that memory pages are moved to the physical memory of the PPU **400** that is accessing the pages more frequently. In an embodiment, the NVLink **410** supports address translation services allowing the PPU **400** to directly access a CPU's page tables and providing full access to CPU memory by the PPU **400**.

(70) In an embodiment, copy engines transfer data between multiple PPUs **400** or between PPUs **400** and CPUs. The copy engines can generate page faults for addresses that are not mapped into the page tables. The memory partition unit **480** can then service the page faults, mapping the addresses into the page table, after which the copy engine can perform the transfer. In a conventional system, memory is pinned (e.g., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing the available memory. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

(71) Data from the memory **404** or other system memory may be fetched by the memory partition unit **480** and stored in the L2 cache **560**, which is located on-chip and is shared between the various processing clusters **450**. As shown, each memory partition unit **480** includes a portion of the L2 cache **560** associated with a corresponding memory device **404**. Lower-level caches may then be implemented in various units within the processing clusters **450**. For example, each of the multiprocessors **540** may implement a level one (L1) cache. The L1 cache is private memory that is dedicated to a particular multiprocessor **540**. Data from the L2 cache **560** may be fetched and stored in each of the L1 caches for processing in the functional units of the Multiprocessors **540**. The L2 cache **560** is coupled to the memory interface **570** and the XBar **470**.

(72) The ROP unit **550** performs graphics raster operations related to pixel color, such as color compression, pixel blending, and the like. The ROP unit **550** also implements depth testing in conjunction with the raster engine **525**, receiving a depth for a sample location associated with a pixel fragment from the culling engine of the raster engine **525**. The depth is tested against a corresponding depth in a depth buffer for a sample location associated with the fragment. If the fragment passes the depth test for the sample location, then the ROP unit **550** updates the depth buffer and transmits a result of the depth test to the raster engine **525**. It will be appreciated that the number of partition units **480** may be different than the number of processing clusters **450** and, therefore, each ROP unit **550** may be coupled to each of the processing clusters **450**. The ROP unit **550** tracks packets received from the different processing clusters **450** and determines which

processing cluster **450** that a result generated by the ROP unit **550** is routed to through the Xbar **470**. Although the ROP unit **550** is included within the memory partition unit **480** in FIG. 5B, in other embodiment, the ROP unit **550** may be outside of the memory partition unit **480**. For example, the ROP unit **550** may reside in the processing cluster **450** or another unit.

(73) FIG. 6A illustrates the multiprocessor **540** of FIG. 5A, in accordance with an embodiment. As shown in FIG. 6A, the multiprocessor **540** includes an instruction cache **605**, one or more scheduler units **610**, a register file **620**, one or more processing cores **650**, one or more special function units (SFUs) **652**, one or more load/store units (LSUs) **654**, an interconnect network **680**, a shared memory/L1 cache **670**.

(74) As described above, the work distribution unit **425** dispatches tasks for execution on the processing clusters **450** of the PPU **400**. The tasks are allocated to a particular DPC **520** within a processing cluster **450** and, if the task is associated with a shader program, the task may be allocated to a multiprocessor **540**. The scheduler unit **610** receives the tasks from the work distribution unit **425** and manages instruction scheduling for one or more thread blocks assigned to the multiprocessor **540**. The scheduler unit **610** schedules thread blocks for execution as warps of parallel threads, where each thread block is allocated at least one warp. In an embodiment, each warp executes 32 threads. The scheduler unit **610** may manage a plurality of different thread blocks, allocating the warps to the different thread blocks and then dispatching instructions from the plurality of different cooperative groups to the various functional units (e.g., cores **650**, SFUs **652**, and LSUs **654**) during each clock cycle.

(75) Cooperative Groups is a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., the `syncthreads()` function). However, programmers would often like to define groups of threads at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

(76) Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

(77) A dispatch unit **615** is configured to transmit instructions to one or more of the functional units. In the embodiment, the scheduler unit **610** includes two dispatch units **615** that enable two different instructions from the same warp to be dispatched during each clock cycle. In alternative embodiments, each scheduler unit **610** may include a single dispatch unit **615** or additional dispatch units **615**.

(78) Each multiprocessor **540** includes a register file **620** that provides a set of registers for the functional units of the multiprocessor **540**. In an embodiment, the register file **620** is divided between each of the functional units such that each functional unit is allocated a dedicated portion of the register file **620**. In another embodiment, the register file **620** is divided between the different warps being executed by the multiprocessor **540**. The register file **620** provides temporary storage for operands connected to the data paths of the functional units.

(79) Each multiprocessor **540** comprises L processing cores **650**. In an embodiment, the multiprocessor **540** includes a large number (e.g., 128, etc.) of distinct processing cores **650**. Each

core **650** may include a fully-pipelined, single-precision, double-precision, and/or mixed precision processing unit that includes a floating-point arithmetic logic unit and an integer arithmetic logic unit. In an embodiment, the floating-point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In an embodiment, the cores **650** include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

(80) Tensor cores configured to perform matrix operations, and, in an embodiment, one or more tensor cores are included in the cores **650**. In particular, the tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In an embodiment, each tensor core operates on a 4×4 matrix and performs a matrix multiply and accumulate operation $D=A \times B+C$, where A, B, C, and D are 4×4 matrices.

(81) In an embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a $4 \times 4 \times 4$ matrix multiply. In practice, Tensor Cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16×16 size matrices spanning all 32 threads of the warp.

(82) Each multiprocessor **540** also comprises M SFUs **652** that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In an embodiment, the SFUs **652** may include a tree traversal unit configured to traverse a hierarchical tree data structure. In an embodiment, the SFUs **652** may include texture unit configured to perform texture map filtering operations. In an embodiment, the texture units are configured to load texture maps (e.g., a 2D array of texels) from the memory **404** and sample the texture maps to produce sampled texture values for use in shader programs executed by the multiprocessor **540**. In an embodiment, the texture maps are stored in the shared memory/L1 cache **570**. The texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In an embodiment, each SM **440** includes two texture units.

(83) Each multiprocessor **540** also comprises N LSUs **654** that implement load and store operations between the shared memory/L1 cache **670** and the register file **620**. Each multiprocessor **540** includes an interconnect network **680** that connects each of the functional units to the register file **620** and the LSU **654** to the register file **620**, shared memory/L1 cache **670**. In an embodiment, the interconnect network **680** is a crossbar that can be configured to connect any of the functional units to any of the registers in the register file **620** and connect the LSUs **654** to the register file and memory locations in shared memory/L1 cache **670**.

(84) The shared memory/L1 cache **670** is an array of on-chip memory that allows for data storage and communication between the multiprocessor **540** and the primitive engine **535** and between threads in the multiprocessor **540**. In an embodiment, the shared memory/L1 cache **670** comprises 128 KB of storage capacity and is in the path from the multiprocessor **540** to the partition unit **480**. The shared memory/L1 cache **670** can be used to cache reads and writes. One or more of the shared memory/L1 cache **670**, L2 cache **560**, and memory **404** are backing stores.

(85) Combining data cache and shared memory functionality into a single memory block provides the best overall performance for both types of memory accesses. The capacity is usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory/L1 cache **670** enables the shared memory/L1 cache **670** to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and

low-latency access to frequently reused data.

(86) When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, the fixed function graphics processing units shown in FIG. 4, are bypassed, creating a much simpler programming model. In the general-purpose parallel computation configuration, the work distribution unit **425** assigns and distributes blocks of threads directly to the DPCs **520**. The threads in a block execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the multiprocessor **540** to execute the program and perform calculations, shared memory/L1 cache **670** to communicate between threads, and the LSU **654** to read and write global memory through the shared memory/L1 cache **670** and the memory partition unit **480**. When configured for general purpose parallel computation, the multiprocessor **540** can also write commands that the scheduler unit **420** can use to launch new work on the DPCs **520**.

(87) The PPU **400** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In an embodiment, the PPU **400** is embodied on a single semiconductor substrate. In another embodiment, the PPU **400** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional PPUs **400**, the memory **404**, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

(88) In an embodiment, the PPU **400** may be included on a graphics card that includes one or more memory devices **404**. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In yet another embodiment, the PPU **400** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

Exemplary Computing System

(89) Systems with multiple GPUs and CPUs are used in a variety of industries as developers expose and leverage more parallelism in applications such as artificial intelligence computing. High-performance GPU-accelerated systems with tens to many thousands of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

(90) FIG. 6B is a conceptual diagram of a processing system **600** implemented using the PPU **400** of FIG. 4, in accordance with an embodiment. The exemplary system **665** may be configured to implement the method **300** shown in FIG. 4. The processing system **600** includes a CPU **630**, switch **612**, and multiple PPUs **400** each and respective memories **404**. The NVLink **410** provides high-speed communication links between each of the PPUs **400**. Although a particular number of NVLink **410** and interconnect **402** connections are illustrated in FIG. 6B, the number of connections to each PPU **400** and the CPU **630** may vary. The switch **612** interfaces between the interconnect **402** and the CPU **630**. The PPUs **400**, memories **404**, and NVLinks **410** may be situated on a single semiconductor platform to form a parallel processing module **625**. In an embodiment, the switch **612** supports two or more protocols to interface between various different connections and/or links.

(91) In another embodiment (not shown), the NVLink **410** provides one or more high-speed communication links between each of the PPUs **400** and the CPU **630** and the switch **612** interfaces between the interconnect **402** and each of the PPUs **400**. The PPUs **400**, memories **404**, and interconnect **402** may be situated on a single semiconductor platform to form a parallel processing module **625**. In yet another embodiment (not shown), the interconnect **402** provides one or more communication links between each of the PPUs **400** and the CPU **630** and the switch **612** interfaces between each of the PPUs **400** using the NVLink **410** to provide one or more high-speed communication links between the PPUs **400**. In another embodiment (not shown), the NVLink **410**

provides one or more high-speed communication links between the PPUs **400** and the CPU **630** through the switch **612**. In yet another embodiment (not shown), the interconnect **402** provides one or more communication links between each of the PPUs **400** directly. One or more of the NVLink **410** high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink **410**.

(92) In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. It should be noted that the term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over utilizing a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the desires of the user. Alternately, the parallel processing module **625** may be implemented as a circuit board substrate and each of the PPUs **400** and/or memories **404** may be packaged devices. In an embodiment, the CPU **630**, switch **612**, and the parallel processing module **625** are situated on a single semiconductor platform.

(93) In an embodiment, the signaling rate of each NVLink **410** is 20 to 25 Gigabits/second and each PPU **400** includes six NVLink **410** interfaces (as shown in FIG. **6B**, five NVLink **410** interfaces are included for each PPU **400**). Each NVLink **410** provides a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 500 Gigabytes/second. The NVLinks **410** can be used exclusively for PPU-to-PPU communication as shown in FIG. **6B**, or some combination of PPU-to-PPU and PPU-to-CPU, when the CPU **630** also includes one or more NVLink **410** interfaces.

(94) In an embodiment, the NVLink **410** allows direct load/store/atomic access from the CPU **630** to each PPU's **400** memory **404**. In an embodiment, the NVLink **410** supports coherency operations, allowing data read from the memories **404** to be stored in the cache hierarchy of the CPU **630**, reducing cache access latency for the CPU **630**. In an embodiment, the NVLink **410** includes support for Address Translation Services (ATS), allowing the PPU **400** to directly access page tables within the CPU **630**. One or more of the NVLinks **410** may also be configured to operate in a low-power mode.

(95) FIG. **6C** illustrates an exemplary system **665** in which the various architecture and/or functionality of the various previous embodiments may be implemented. The exemplary system **665** may be configured to implement the method **300** shown in FIG. **4**.

(96) As shown, a system **665** is provided including at least one central processing unit **630** that is connected to a communication bus **675**. The communication bus **675** may be implemented using any suitable protocol, such as PCI (Peripheral Component Interconnect), PCI-Express, AGP (Accelerated Graphics Port), HyperTransport, or any other bus or point-to-point communication protocol(s). The system **665** also includes a main memory **640**. Control logic (software) and data are stored in the main memory **640** which may take the form of random-access memory (RAM).

(97) The system **665** also includes input devices **660**, the parallel processing system **625**, and display devices **645**, e.g., a conventional CRT (cathode ray tube), LCD (liquid crystal display), LED (light emitting diode), plasma display or the like. User input may be received from the input devices **660**, e.g., keyboard, mouse, touchpad, microphone, and the like. Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the system **665**. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

(98) Further, the system **665** may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface **635** for communication purposes.

(99) The system **665** may also include a secondary storage (not shown). The secondary storage

include, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive reads from and/or writes to a removable storage unit in a well-known manner.

(100) Computer programs, or computer control logic algorithms, may be stored in the main memory **640** and/or the secondary storage. Such computer programs, when executed, enable the system **665** to perform various functions. The memory **640**, the storage, and/or any other storage are possible examples of computer-readable media.

(101) The architecture and/or functionality of the various previous figures may be implemented in the context of a general computer system, a circuit board system, a game console system dedicated for entertainment purposes, an application-specific system, and/or any other desired system. For example, the system **665** may take the form of a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, a mobile phone device, a television, workstation, game consoles, embedded system, and/or any other type of logic.

(102) While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary embodiments, but should be defined only in accordance with the following claims and their equivalents.

Graphics Processing Pipeline

(103) In an embodiment, the PPU **400** comprises a graphics processing unit (GPU). The PPU **400** is configured to receive commands that specify shader programs for processing graphics data.

Graphics data may be defined as a set of primitives such as points, lines, triangles, quads, triangle strips, and the like. Typically, a primitive includes data that specifies a number of vertices for the primitive (e.g., in a model-space coordinate system) as well as attributes associated with each vertex of the primitive. The PPU **400** can be configured to process the graphics primitives to generate a frame buffer (e.g., pixel data for each of the pixels of the display).

(104) An application writes model data for a scene (e.g., a collection of vertices and attributes) to a memory such as a system memory or memory **404**. The model data defines each of the objects that may be visible on a display. The application then makes an API call to the driver kernel that requests the model data to be rendered and displayed. The driver kernel reads the model data and writes commands to the one or more streams to perform operations to process the model data. The commands may reference different shader programs to be implemented on the multiprocessors **540** of the PPU **400** including one or more of a vertex shader, hull shader, domain shader, geometry shader, and a pixel shader. For example, one or more of the multiprocessors **540** may be configured to execute a vertex shader program that processes a number of vertices defined by the model data. In an embodiment, the different multiprocessors **540** may be configured to execute different shader programs concurrently. For example, a first subset of multiprocessors **540** may be configured to execute a vertex shader program while a second subset of multiprocessors **540** may be configured to execute a pixel shader program. The first subset of multiprocessors **540** processes vertex data to produce processed vertex data and writes the processed vertex data to the L2 cache **460** and/or the memory **404**. After the processed vertex data is rasterized (e.g., transformed from three-dimensional data into two-dimensional data in screen space) to produce fragment data, the second subset of multiprocessors **540** executes a pixel shader to produce processed fragment data, which is then blended with other processed fragment data and written to the frame buffer in memory **404**. The vertex shader program and pixel shader program may execute concurrently, processing different data from the same scene in a pipelined fashion until all of the model data for the scene has been rendered to the frame buffer. Then, the contents of the frame buffer are transmitted to a

display controller for display on a display device.

(105) FIG. 7 is a conceptual diagram of a graphics processing pipeline **700** implemented by the PPU **400** of FIG. 4, in accordance with an embodiment. The graphics processing pipeline **700** is an abstract flow diagram of the processing steps implemented to generate 2D computer-generated images from 3D geometry data. As is well-known, pipeline architectures may perform long latency operations more efficiently by splitting up the operation into a plurality of stages, where the output of each stage is coupled to the input of the next successive stage. Thus, the graphics processing pipeline **700** receives input data **701** that is transmitted from one stage to the next stage of the graphics processing pipeline **700** to generate output data **702**. In an embodiment, the graphics processing pipeline **700** may represent a graphics processing pipeline defined by the OpenGL® API. As an option, the graphics processing pipeline **700** may be implemented in the context of the functionality and architecture of the previous Figures and/or any subsequent Figure(s).

(106) As shown in FIG. 7, the graphics processing pipeline **700** comprises a pipeline architecture that includes a number of stages. The stages include, but are not limited to, a data assembly stage **710**, a vertex shading stage **720**, a primitive assembly stage **730**, a geometry shading stage **740**, a viewport scale, cull, and clip (VSCC) stage **750**, a rasterization stage **760**, a fragment shading stage **770**, and a raster operations stage **780**. In an embodiment, the input data **701** comprises commands that configure the processing units to implement the stages of the graphics processing pipeline **700** and geometric primitives (e.g., points, lines, triangles, quads, triangle strips or fans, etc.) to be processed by the stages. The output data **702** may comprise pixel data (e.g., color data) that is copied into a frame buffer or other type of surface data structure in a memory.

(107) The data assembly stage **710** receives the input data **701** that specifies vertex data for high-order surfaces, primitives, or the like. The data assembly stage **710** collects the vertex data in a temporary storage or queue, such as by receiving a command from the host processor that includes a pointer to a buffer in memory and reading the vertex data from the buffer. The vertex data is then transmitted to the vertex shading stage **720** for processing.

(108) The vertex shading stage **720** processes vertex data by performing a set of operations (e.g., a vertex shader or a program) once for each of the vertices. Vertices may be, e.g., specified as a 4-coordinate vector (e.g., $\langle x, y, z, w \rangle$) associated with one or more vertex attributes (e.g., color, texture coordinates, surface normal, etc.). The vertex shading stage **720** may manipulate individual vertex attributes such as position, color, texture coordinates, and the like. In other words, the vertex shading stage **720** performs operations on the vertex coordinates or other vertex attributes associated with a vertex. Such operations commonly including lighting operations (e.g., modifying color attributes for a vertex) and transformation operations (e.g., modifying the coordinate space for a vertex). For example, vertices may be specified using coordinates in an object-coordinate space, which are transformed by multiplying the coordinates by a matrix that translates the coordinates from the object-coordinate space into a world space or a normalized-device-coordinate (NDC) space. The vertex shading stage **720** generates transformed vertex data that is transmitted to the primitive assembly stage **730**.

(109) The primitive assembly stage **730** collects vertices output by the vertex shading stage **720** and groups the vertices into geometric primitives for processing by the geometry shading stage **740**. For example, the primitive assembly stage **730** may be configured to group every three consecutive vertices as a geometric primitive (e.g., a triangle) for transmission to the geometry shading stage **740**. In some embodiments, specific vertices may be reused for consecutive geometric primitives (e.g., two consecutive triangles in a triangle strip may share two vertices). The primitive assembly stage **730** transmits geometric primitives (e.g., a collection of associated vertices) to the geometry shading stage **740**.

(110) The geometry shading stage **740** processes geometric primitives by performing a set of operations (e.g., a geometry shader or program) on the geometric primitives. Tessellation operations may generate one or more geometric primitives from each geometric primitive. In other

words, the geometry shading stage **740** may subdivide each geometric primitive into a finer mesh of two or more geometric primitives for processing by the rest of the graphics processing pipeline **700**. The geometry shading stage **740** transmits geometric primitives to the viewport SCC stage **750**.

(111) In an embodiment, the graphics processing pipeline **700** may operate within a multiprocessor and the vertex shading stage **720**, the primitive assembly stage **730**, the geometry shading stage **740**, the fragment shading stage **770**, and/or hardware/software associated therewith, may sequentially perform processing operations. Once the sequential processing operations are complete, in an embodiment, the viewport SCC stage **750** may utilize the data. In an embodiment, primitive data processed by one or more of the stages in the graphics processing pipeline **700** may be written to a cache (e.g., L1 cache, a vertex cache, etc.). In this case, in an embodiment, the viewport SCC stage **750** may access the data in the cache. In an embodiment, the viewport SCC stage **750** and the rasterization stage **760** are implemented as fixed function circuitry.

(112) The viewport SCC stage **750** performs viewport scaling, culling, and clipping of the geometric primitives. Each surface being rendered to is associated with an abstract camera position. The camera position represents a location of a viewer looking at the scene and defines a viewing frustum that encloses the objects of the scene. The viewing frustum may include a viewing plane, a rear plane, and four clipping planes. Any geometric primitive entirely outside of the viewing frustum may be culled (e.g., discarded) because the geometric primitive will not contribute to the final rendered scene. Any geometric primitive that is partially inside the viewing frustum and partially outside the viewing frustum may be clipped (e.g., transformed into a new geometric primitive that is enclosed within the viewing frustum). Furthermore, geometric primitives may each be scaled based on a depth of the viewing frustum. All potentially visible geometric primitives are then transmitted to the rasterization stage **760**.

(113) The rasterization stage **760** converts the 3D geometric primitives into 2D fragments (e.g., capable of being utilized for display, etc.). The rasterization stage **760** may be configured to utilize the vertices of the geometric primitives to set up a set of plane equations from which various attributes can be interpolated. The rasterization stage **760** may also compute a coverage mask for a plurality of pixels that indicates whether one or more sample locations for the pixel intercept the geometric primitive. In an embodiment, z-testing may also be performed to determine if the geometric primitive is occluded by other geometric primitives that have already been rasterized. The rasterization stage **760** generates fragment data (e.g., interpolated vertex attributes associated with a particular sample location for each covered pixel) that are transmitted to the fragment shading stage **770**.

(114) The fragment shading stage **770** processes fragment data by performing a set of operations (e.g., a fragment shader or a program) on each of the fragments. The fragment shading stage **770** may generate pixel data (e.g., color values) for the fragment such as by performing lighting operations or sampling texture maps using interpolated texture coordinates for the fragment. The fragment shading stage **770** generates pixel data that is transmitted to the raster operations stage **780**.

(115) The raster operations stage **780** may perform various operations on the pixel data such as performing alpha tests, stencil tests, and blending the pixel data with other pixel data corresponding to other fragments associated with the pixel. When the raster operations stage **780** has finished processing the pixel data (e.g., the output data **702**), the pixel data may be written to a render target such as a frame buffer, a color buffer, or the like.

(116) It will be appreciated that one or more additional stages may be included in the graphics processing pipeline **700** in addition to or in lieu of one or more of the stages described above. Various implementations of the abstract graphics processing pipeline may implement different stages. Furthermore, one or more of the stages described above may be excluded from the graphics processing pipeline in some embodiments (such as the geometry shading stage **740**). Other types of

graphics processing pipelines are contemplated as being within the scope of the present disclosure. Furthermore, any of the stages of the graphics processing pipeline **700** may be implemented by one or more dedicated hardware units within a graphics processor such as PPU **400**. Other stages of the graphics processing pipeline **700** may be implemented by programmable hardware units such as the multiprocessors **540** of the PPU **400**.

(117) The graphics processing pipeline **700** may be implemented via an application executed by a host processor, such as a CPU. In an embodiment, a device driver may implement an application programming interface (API) that defines various functions that can be utilized by an application in order to generate graphical data for display. The device driver is a software program that includes a plurality of instructions that control the operation of the PPU **400**. The API provides an abstraction for a programmer that lets a programmer utilize specialized graphics hardware, such as the PPU **400**, to generate the graphical data without requiring the programmer to utilize the specific instruction set for the PPU **400**. The application may include an API call that is routed to the device driver for the PPU **400**. The device driver interprets the API call and performs various operations to respond to the API call. In some instances, the device driver may perform operations by executing instructions on the CPU. In other instances, the device driver may perform operations, at least in part, by launching operations on the PPU **400** utilizing an input/output interface between the CPU and the PPU **400**. In an embodiment, the device driver is configured to implement the graphics processing pipeline **700** utilizing the hardware of the PPU **400**.

(118) Various programs may be executed within the PPU **400** in order to implement the various stages of the graphics processing pipeline **700**. For example, the device driver may launch a kernel on the PPU **400** to perform the vertex shading stage **720** on one multiprocessors **540** (or multiple multiprocessors **540**). The device driver (or the initial kernel executed by the PPU **400**) may also launch other kernels on the PPU **400** to perform other stages of the graphics processing pipeline **700**, such as the geometry shading stage **740** and the fragment shading stage **770**. In addition, some of the stages of the graphics processing pipeline **700** may be implemented on fixed unit hardware such as a rasterizer or a data assembler implemented within the PPU **400**. It will be appreciated that results from one kernel may be processed by one or more intervening fixed function hardware units before being processed by a subsequent kernel on an multiprocessors **540**.

Machine Learning

(119) Deep neural networks (DNNs) developed on processors, such as the PPU **400** have been used for diverse use cases, from self-driving cars to faster drug development, from automatic image captioning in online image databases to smart real-time language translation in video chat applications. Deep learning is a technique that models the neural learning process of the human brain, continually learning, continually getting smarter, and delivering more accurate results more quickly over time. A child is initially taught by an adult to correctly identify and classify various shapes, eventually being able to identify shapes without any coaching. Similarly, a deep learning or neural learning system needs to be trained in object recognition and classification for it get smarter and more efficient at identifying basic objects, occluded objects, etc., while also assigning context to objects.

(120) At the simplest level, neurons in the human brain look at various inputs that are received, importance levels are assigned to each of these inputs, and output is passed on to other neurons to act upon. An artificial neuron or perceptron is the most basic model of a neural network. In one example, a perceptron may receive one or more inputs that represent various features of an object that the perceptron is being trained to recognize and classify, and each of these features is assigned a certain weight based on the importance of that feature in defining the shape of an object.

(121) A deep neural network (DNN) model includes multiple layers of many connected nodes (e.g., perceptrons, Boltzmann machines, radial basis functions, convolutional layers, etc.) that can be trained with enormous amounts of input data to quickly solve complex problems with high accuracy. In one example, a first layer of the DNN model breaks down an input image of an

automobile into various sections and looks for basic patterns such as lines and angles. The second layer assembles the lines to look for higher level patterns such as wheels, windshields, and mirrors. The next layer identifies the type of vehicle, and the final few layers generate a label for the input image, identifying the model of a specific automobile brand.

(122) Once the DNN is trained, the DNN can be deployed and used to identify and classify objects or patterns in a process known as inference. Examples of inference (the process through which a DNN extracts useful information from a given input) include identifying handwritten numbers on checks deposited into ATM machines, identifying images of friends in photos, delivering movie recommendations to over fifty million users, identifying and classifying different types of automobiles, pedestrians, and road hazards in driverless cars, or translating human speech in real-time.

(123) During training, data flows through the DNN in a forward propagation phase until a prediction is produced that indicates a label corresponding to the input. If the neural network does not correctly label the input, then errors between the correct label and the predicted label are analyzed, and the weights are adjusted for each feature during a backward propagation phase until the DNN correctly labels the input and other inputs in a training dataset. Training complex neural networks requires massive amounts of parallel computing performance, including floating-point multiplications and additions that are supported by the PPU **400**. Inferencing is less compute-intensive than training, being a latency-sensitive process where a trained neural network is applied to new inputs it has not seen before to classify images, translate speech, and generally infer new information.

(124) Neural networks rely heavily on matrix math operations, and complex multi-layered networks require tremendous amounts of floating-point performance and bandwidth for both efficiency and speed. With thousands of processing cores, optimized for matrix math operations, and delivering tens to hundreds of TFLOPS of performance, the PPU **400** is a computing platform capable of delivering performance required for deep neural network-based artificial intelligence and machine learning applications.

Claims

1. A system comprising: a processing device including a first cache; a system memory, operatively coupled with the processing device; and a parallel processing unit (PPU), operatively coupled with the processing device, wherein the PPU includes a second cache, and wherein the PPU is to: receive a cache operation to modify cache lines of the second cache associated with a first aperture of the system memory; and identify a first subset of cache lines of the second cache, wherein the first subset of cache lines is associated with the first aperture of the system memory; identify a second subset of cache lines of the second cache, wherein the second subset of cache lines is associated with a second aperture of the system memory; and modify the first subset of cache lines as specified by the cache operation.
2. The system of claim 1, wherein the cache operation is an invalidate operation, and wherein to modify the first subset of cache lines, the PPU is to: write data stored at the first subset of cache lines back to the first cache; and invalidate the first subset of cache lines.
3. The system of claim 1, wherein the cache operation is a flush operation, and wherein to modify the first subset of cache lines, PPU is to: write data stored at the first subset of cache lines back to the first cache; and maintain a clean state of the first subset of cache lines.
4. The system of claim 1, wherein the first aperture is a non-coherent aperture of the system memory, and the second aperture is a coherent aperture of the system memory.
5. The system of claim 1, wherein the PPU is further to: cause the second subset of cache lines to be maintained within the second cache.
6. The system of claim 1, wherein the first subset of cache lines and the second subset of cache

lines are identified based on identifiers associated with a non-coherent system memory aperture and a coherent system memory aperture, respectively.

7. The system of claim 1, wherein the PPU and the processing device are interconnected via an interface using a common hardware interface (CHI) protocol.

8. The system of claim 1, wherein the first subset of cache lines can further be differentiated and invalidated based on a process identifier that indicates one of a plurality of processes associated with the processing device.

9. The system of claim 1, wherein coherency of the second subset of cache lines is managed by hardware associated with the processing device.

10. A method comprising: receiving, at a parallel processing unit (PPU) including a first cache, a cache operation to modify cache lines of the first cache associated with a first aperture of a system memory of a processing device; identifying a first subset of cache lines of the first cache, wherein the first subset of cache lines is associated with the first aperture of the system memory; identifying a second subset of cache lines of the first cache, wherein the second subset of cache lines is associated with a second aperture of the system memory; and modifying the first subset of cache lines as specified by the cache operation.

11. The method of claim 10, wherein the cache operation is an invalidate operation, and wherein modifying the first subset of cache lines comprises: writing data stored at the first subset of cache lines back to a second cache of the processing device; and invalidating the first subset of cache lines.

12. The method of claim 10, wherein the cache operation is a flush operation, and wherein modifying the first subset of cache lines comprises: writing data stored at the first subset of cache lines back to a second cache of the processing device; and maintaining a clean state of the first subset of cache lines.

13. The method of claim 10, wherein the first aperture is a non-coherent aperture of the system memory, and the second aperture is a coherent aperture of the system memory.

14. The method of claim 10, further comprising: causing the second subset of cache lines to be maintained within the first cache.

15. The method of claim 10, wherein the first subset of cache lines and the second subset of cache lines are identified based on identifiers associated with a non-coherent aperture and a coherent aperture, respectively.

16. The method of claim 10, wherein the PPU and the processing device are interconnected via an interface using a common hardware interface (CHI) protocol.

17. The method of claim 10, wherein the first subset of cache lines can further be differentiated and invalidated based on a process identifier that indicates one of a plurality of processes associated with the processing device.

18. The method of claim 10, wherein coherency of the second subset of cache lines is managed by hardware associated with the processing device.

19. One or more processors comprising processing circuitry to: receive, at a parallel processing unit (PPU) including a first cache, a cache operation to modify cache lines of the first cache associated with a first aperture of a system memory of a processing device; identify a first subset of cache lines of the first cache, wherein the first subset of cache lines is associated with the first aperture of the system memory; identify a second subset of cache lines of the first cache, wherein the second subset of cache lines is associated with a second aperture of the system memory; and modify the first subset of cache lines as specified by the cache operation.

20. The one or more processors of claim 19, wherein the cache operation is an invalidate operation, and wherein to modify the first subset of cache lines, the processing circuitry is to: write data stored at the first subset of cache lines back to a second cache of the processing device; and invalidate the first subset of cache lines as invalid.
