

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent

12417183

Kind Code

B2

Date of Patent

September 16, 2025

Inventor(s)

Halaharivi; Raja V. S. et al.

Page request interface (PRI) handler operational support in relation to caching host memory address translation data in a memory sub-system

Abstract

A processing device includes host interface circuitry to interact with a host system and an address translation circuit to handle address translation requests to the host system from host interface circuits. The address translation circuit includes a cache to store address translations associated with the address translation requests for future access by host interface circuits. A page request interface (PRI) handler tracks translation miss messages received from the host interface circuits, each translation miss message including a virtual address of a miss at the cache. The PRI handler removes duplicate translation miss messages having an identical virtual address and creates page miss requests from non-duplicate translation miss messages that are categorized into page request groups, each page request group corresponding to a host interface circuit of the host interface circuits. The PRI handler queues the page request groups to be sent to a translation agent of the host system.

Inventors: Halaharivi; Raja V. S. (Gilroy, CA), Sharma; Prateek (San Jose, CA), Chakrabarty; Sumangal (Campbell, CA), Gaddam; Venkat R. (Fremont, CA)

Applicant: Micron Technology, Inc. (Boise, ID)

Family ID: 1000008820237

Assignee: Micron Technology, Inc. (Boise, ID)

Appl. No.: 18/505302

Filed: November 09, 2023

Prior Publication Data

Document Identifier

Publication Date

US 20240160577 A1

May. 16, 2024

Related U.S. Application Data

Publication Classification

Int. Cl.: G06F12/1009 (20160101); G06F12/123 (20160101)

U.S. Cl.:

CPC G06F12/1009 (20130101); G06F12/123 (20130101); G06F2212/1021 (20130101)

Field of Classification Search

CPC: G06F (12/1009); G06F (12/123); G06F (2212/1021)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2011/0138149	12/2010	Karlsson	711/E12.001	G06F 12/1027
2011/0271057	12/2010	Karlsson	711/E12.024	G06F 12/0895
2019/0179757	12/2018	Walker	N/A	G06F 13/161
2019/0182168	12/2018	Bastide	N/A	H04L 47/24
2022/0043753	12/2021	Cohen	N/A	G06F 12/0871

Primary Examiner: Mackall; Larry T

Attorney, Agent or Firm: Lowenstein Sandler LLP

Background/Summary

REFERENCE TO RELATED APPLICATION (1) This application claims the benefit of U.S. Provisional Patent Application No. 63/424,409, filed Nov. 10, 2022, the entirety of which is incorporated herein by reference.

TECHNICAL FIELD

(1) The present disclosure generally relates to a memory system, and more specifically, relates to page request interface support in caching host memory address translation data in a memory sub-system.

BACKGROUND

(2) A memory sub-system can include one or more memory components that store data. The memory components can be, for example, non-volatile memory components and volatile memory components. In general, a host system can utilize a memory sub-system to store data at the memory components and to retrieve data from the memory components.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) The present disclosure will be understood more fully from the detailed description given below and from the accompanying drawings of various embodiments of the disclosure.
- (2) FIG. 1 illustrates an example computing environment for page request interface support in caching host memory address translation data in a memory sub-system in accordance with some embodiments of the present disclosure.
- (3) FIG. 2 is a schematic block diagram of a system (or device) implementing peripheral component interface express (PCIe) and non-volatile memory express (NVMe) functionality within which the disclosed page request interface support operates in accordance with some embodiments.
- (4) FIG. 3 is a schematic diagram illustrating an embodiment of a controller memory buffer (CMB) manager of a storage device.
- (5) FIG. 4 is an example physical host interface between a host system and a memory sub-system implementing caching host memory address translation data in accordance with some embodiments.
- (6) FIG. 5 is a system in which a memory sub-system controller contains an address translation cache (ATC) in accordance with some embodiments.
- (7) FIG. 6 is a memory sub-system for page request interface support in caching host memory address translation data for multiple host interface circuits in accordance with some embodiments.
- (8) FIG. 7 is an example memory sub-system controller including an address translation circuit implementing caching host memory address translation data in accordance with some embodiments.
- (9) FIG. 8 is a flow chart of an example method of caching host memory address translation data in a memory sub-system in accordance with some embodiments.
- (10) FIG. 9 is a block diagram of a system for host interface circuits interacting with a page request interface handler in accordance with some embodiments.
- (11) FIG. 10 is a schematic block diagram of a system including a processing bridge circuit implementing page request interface support in caching host memory address translation data in a memory sub-system in accordance with some embodiments.
- (12) FIG. 11 is a flow chart of an example method of page request interface support in caching host memory address translation data in a memory sub-system in accordance with some embodiments.
- (13) FIG. 12 is a block diagram of an example computer system in which embodiments of the present disclosure can operate.

DETAILED DESCRIPTION

- (14) Aspects of the present disclosure are directed to page request interface support in caching host memory address translation data in a memory sub-system. A memory sub-system can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of storage devices and memory modules are described below in conjunction with FIG. 1. In general, a host system can utilize a memory sub-system that includes one or more memory components (also hereinafter referred to as “memory devices”). The host system can provide data to be stored at the memory sub-system and can request data to be retrieved from the memory sub-system.
- (15) In requesting data be written to or read from a memory device, the host system typically generates memory commands (e.g., an erase (or unmap) command, a write command, or a read command) that are sent to a memory sub-system controller (e.g., processing device or “controller”). The controller then executes on these memory commands to perform an erase (or unmap) operation, a write operation, or a read operation at the memory device. Because the host operates in logical addresses, which are referred to as virtual addresses (or guest physical addresses) in the context of virtual machines (VMs) that run on the host system, the host system includes a root complex that serves as a connection between the physical and virtual components of the host system and a peripheral control interconnect express (PCIe) bus. This PCIe root complex can

generate transaction requests (to include address translation) requests on behalf of entities of the host system, such as a virtual processing device in one of the VMs.

(16) The host system typically further includes a translation agent (TA) that performs translations, on behalf of the controller, of virtual addresses to physical addresses. To do so, the TA is configured to communicate with translation requests/responses through the PCIe root complex. In some systems, the TA is also known as an input/output memory management unit (IOMMU) that is executed by a hypervisor or virtual machine manager running on the host system. Thus, the TA can be a hardware component or software (IOMMU) with a dedicated driver.

(17) The controller in these systems can be configured to include an address translation circuit, more specifically referred to as an address translation service (ATS), that is to request the TA to perform certain address translations from a virtual (or logical) address to an available (or assigned) physical address of the memory device. In this way, the address translation circuit (or ATS) dynamically determines address translations based on the virtual address located in a corresponding memory command that is queued within host memory. Different aspects of the ATS obviate the need to pin a substantial amount of memory associated with an application being run by the host system.

(18) Especially in support of multiple non-volatile memory express (NVMe) devices, the need to continually request the TA to perform address translations is a bottleneck and affects performance in terms of speed, latency, and quality-of-service in fulfilling memory commands. Performance can be increasingly impacted as submission, completion, I/O, and administrative queues located within the host memory get larger and the speeds of media of the memory devices increase. For example, the number of address translation requests and responses for command queues as well as for direct memory access (DMA) addresses can be slowed by having to move back and forth across the PCIe bus, which also generates additional I/O traffic that slows the entire memory sub-system.

(19) Aspects of the present disclosure address the above and other deficiencies by implementing, within the address translation circuit of host interface circuitry within the controller, an address translation cache (ATC) that stores address translations corresponding to incoming address translation requests from host interface (HIF) circuits of the host interface circuitry. The ATC can store the address translations, associated with the address translation requests, for future access by the host interface circuits. These address translation requests, for example, may be related to processing of memory commands as well as the handling of DMA operations or commands. In this way, when a cached address translation matches a subsequent (or later) address translation request from a HIF circuit (e.g., hits at the cache), the address translation circuit can retrieve and return the cached address translation to the HIF circuit without having to request the TA to perform the translation on behalf of the controller.

(20) In some embodiments, for each memory command within a submission queue of the host memory, the address translation circuit can store a first address translation in the ATC corresponding to a current page targeted by the memory command (referenced in an address translation request) and store a second address translation in the ATC for a subsequent page that sequentially follows the current page according to virtual address numbering. This look-ahead buffering in the ATC of address translations for a predetermined number of submission queues enables greatly reducing the number of misses at the ATC while keeping a size of the ATC reasonable given the expense of cache memory, e.g., static random access memory (SRAM), availability at the controller. The hit rate at the cache can further be increased by this approach when the command (and other) queues in the host memory are arranged to sequentially store memory commands according to virtual addresses.

(21) In most devices or systems, DMA operations cannot handle page faults, e.g., there is no way to resolve the page faults (such as via paging) to successfully perform a given DMA command. To prevent page faults with DMA operations, the host system pins pages in host memory (whether source or destination of a DMA) before a driver submits I/O data to the device. Pinning a page

within the host address space of memory guarantees immediate execution of a DMA operation when triggered by an interrupt. If an I/O device supports a large amount of pre-allocated queues and data buffers, a significant chunk of host memory is reserved upfront even if the I/O device is not using all this pinned memory at the same time. For most applications and workload situations, all that pinned memory typically sits idle.

(22) Thus, while the address translation circuit would seem to enable the ability to forego pinning of memory (by accessing the ATC), there are instances of missing at the ATC (or cache of the ATS), which would generate the need for paging memory. But, because DMA operations do not support page faults, the address translation circuit may still be an incomplete solution for DMA operations. In order to resolve this problem, in some embodiments, an additional page interface request (PRI) handler is added to the host interface design that automates, from a memory controller perspective, sending PRI-related page miss requests to the host system to request the host system re-pin a physical page of memory to a virtual address of each respective page miss request, thus making these pages available again in host memory.

(23) In some embodiments, for example, the PRI handler tracks translation miss messages received from the host interface circuits, each translation miss message including a virtual address of a miss at the cache. The PRI handler can further remove duplicate translation miss messages having an identical virtual address, e.g., due overlap in virtual addresses concurrently coming from multiple host interface circuits. The PM handler can further create page miss requests from non-duplicate translation miss messages that are categorized into page request groups. For example, a page request group corresponds to one of the host interface circuits.

(24) In these embodiments, the PRI handler further queues the page request groups to be sent to a translation agent of the host system, e.g., for handling of the page miss requests and sending back confirmations in the form of page request responses. In some embodiments, the page request responses trigger the PRI handler to send out restart messages to the host interface circuits along with identification of which hardware pipeline portion or queue that can be restarted now that the page is available in memory. In these embodiments, the address translations to pages that have been re-pinned can further be stored in the ATC for future accesses by any of the HIF circuits.

(25) Therefore, advantages of the systems and methods implemented in accordance with some embodiments of the present disclosure include, but are not limited to, improving performance of the memory sub-system in terms of speed, latency, and throughput of handling memory commands. Part of the reason for increased performance is reducing the I/O traffic over the PCIe buses of the memory sub-system and at the host TA. The disclosed address translation circuit can also reduce the likelihood that previously cached translations will be invalidated and have to be re-fetched from the TA of the host system. Additional advantages of the PRI handler involvement in the caching of host memory address translation data in the memory sub-system include avoiding the need to pin large amounts of host memory to support DMA operations or the like by requesting pages be re-pinned when needed. Other advantages will be apparent to those skilled in the art of address translations within memory sub-systems, which will be discussed hereinafter. Additional details of these techniques are provided below with respect to FIGS. 1-11.

(26) FIG. 1 illustrates an example computing environment **100** that includes a memory sub-system **110** in accordance with some embodiments of the present disclosure. The memory sub-system **110** can include media, such as one or more volatile memory devices (e.g., memory device **140**), one or more non-volatile memory devices (e.g., memory device **130**), or a combination of such.

(27) A memory sub-system **110** can be a storage device, a memory module, or a hybrid of a storage device and memory module. Examples of a storage device include a solid-state drive (SSD), a flash drive, a universal serial bus (USB) flash drive, an embedded Multi-Media Controller (eMMC) drive, a Universal Flash Storage (UFS) drive, and a hard disk drive (HDD). Examples of memory modules include a dual in-line memory module (DIMM), a small outline DIMM (SO-DIMM), and a non-volatile dual in-line memory module (NVDIMM).

(28) The computing environment **100** can include a host system **120** that is coupled to one or more memory sub-systems **110**. In some embodiments, the host system **120** is coupled to different types of memory sub-system **110**. FIG. **1** illustrates one example of a host system **120** coupled to one memory sub-system **110**. The host system **120** uses the memory sub-system **110**, for example, to write data to the memory sub-system **110** and read data from the memory sub-system **110**. As used herein, “coupled to” generally refers to a connection between components, which can be an indirect communicative connection or direct communicative connection (e.g., without intervening components), whether wired or wireless, including connections such as electrical, optical, magnetic, etc.

(29) The host system **120** can be a computing device such as a desktop computer, laptop computer, network server, mobile device, embedded computer (e.g., one included in a vehicle, industrial equipment, or a networked commercial device), or such computing device that includes a memory and a processing device. The host system **120** can be coupled to the memory sub-system **110** via a physical host interface. Examples of a physical host interface include, but are not limited to, a serial advanced technology attachment (SATA) interface, a peripheral component interconnect express (PCIe) interface, universal serial bus (USB) interface, Fibre Channel, Serial Attached SCSI (SAS), etc. The physical host interface can be used to transmit data between the host system **120** and the memory sub-system **110**. The host system **120** can further utilize an NVM Express (NVMe) interface to access the memory components (e.g., memory devices **130**) when the memory sub-system **110** is coupled with the host system **120** by the PCIe interface. The physical host interface can provide an interface for passing control, address, data, and other signals between the memory sub-system **110** and the host system **120**.

(30) The memory devices can include any combination of the different types of non-volatile memory devices and/or volatile memory devices. The volatile memory devices (e.g., memory device **140**) can be, but are not limited to, random access memory (RAM), such as dynamic random access memory (DRAM) and synchronous dynamic random access memory (SDRAM).

(31) Some examples of non-volatile memory devices (e.g., memory device **130**) include NOT-AND (NAND) type flash memory and write-in-place memory, such as three-dimensional cross-point (“3D cross-point”) memory. A 3D cross-point memory device is a cross-point array of non-volatile memory cells that can perform bit storage based on a change of bulk resistance, in conjunction with a stackable cross-gridded data access array. Additionally, in contrast to many flash-based memories, cross-point non-volatile memory can perform a write-in-place operation, where a non-volatile memory cell can be programmed without the non-volatile memory cell being previously erased.

(32) Each of the memory devices **130** can include one or more arrays of memory cells such as single level cells (SLCs), multi-level cells (MLCs), triple level cells (TLCs), or quad-level cells (QLCs). In some embodiments, a particular memory component can include an SLC portion, and an MLC portion, a TLC portion, or a QLC portion of memory cells. Each of the memory cells can store one or more bits of data used by the host system **120**. Furthermore, the memory cells of the memory devices **130** can be grouped to form pages that can refer to a unit of the memory component used to store data. With some types of memory (e.g., NAND), pages can be grouped to form blocks. Some types of memory, such as 3D cross-point, can group pages across die and channels to form management units (MUs).

(33) Although non-volatile memory components such as NAND type flash memory and 3D cross-point are described, the memory device **130** can be based on any other type of non-volatile memory, such as read-only memory (ROM), phase change memory (PCM), magneto random access memory (MRAM), NOT-OR (NOR) flash memory, electrically erasable programmable read-only memory (EEPROM).

(34) The memory sub-system controller **115** can communicate with the memory devices **130** to perform operations such as reading data, writing data, or erasing data at the memory devices **130**

and other such operations. The memory sub-system controller **115** can include hardware such as one or more integrated circuits and/or discrete components, a buffer memory, or a combination thereof. The hardware can include a digital circuitry with dedicated (i.e., hard-coded) logic to perform the operations described herein. The memory sub-system controller **115** can be a microcontroller, special purpose logic circuitry (e.g., a field programmable gate array (FPGA), an application specific integrated circuit (ASIC), etc.), or other suitable processor.

(35) The memory sub-system controller **115** can include a processor (processing device) **117** configured to execute instructions stored in local memory **119**. In the illustrated example, the local memory **119** of the memory sub-system controller **115** includes an embedded memory configured to store instructions for performing various processes, operations, logic flows, and routines that control operation of the memory sub-system **110**, including handling communications between the memory sub-system **110** and the host system **120**.

(36) In some embodiments, the local memory **119** can include memory registers storing memory pointers, fetched data, etc. The local memory **119** can also include read-only memory (ROM) for storing micro-code. While the example memory sub-system **110** in FIG. **1** has been illustrated as including the memory sub-system controller **115**, in another embodiment of the present disclosure, a memory sub-system **110** may not include a memory sub-system controller **115**, and may instead rely upon external control (e.g., provided by an external host, or by a processor or controller separate from the memory sub-system).

(37) In general, the memory sub-system controller **115** can receive commands or operations from the host system **120** and can convert the commands or operations into instructions or appropriate commands to achieve the desired access to the memory devices **130**. The memory sub-system controller **115** can be responsible for other operations such as wear leveling operations, garbage collection operations, error detection and error-correcting code (ECC) operations, encryption operations, caching operations, and address translations between a logical block address and a physical block address that are associated with the memory devices **130**. The memory sub-system controller **115** can further include host interface circuitry to communicate with the host system **120** via the physical host interface. The host interface circuitry can convert the commands received from the host system into command instructions to access the memory devices **130** as well as convert responses associated with the memory devices **130** into information for the host system **120**.

(38) The memory sub-system **110** can also include additional circuitry or components that are not illustrated. In some embodiments, the memory sub-system **110** can include a cache or buffer (e.g., DRAM) and address circuitry (e.g., a row decoder and a column decoder) that can receive an address from the memory sub-system controller **115** and decode the address to access the memory devices **130**.

(39) In some embodiments, the memory devices **130** include local media controllers **135** that operate in conjunction with memory sub-system controller **115** to execute operations on one or more memory cells of the memory devices **130**. In some embodiments, the memory devices **130** are managed memory devices, which is a raw memory device combined with a local controller (e.g., local media controller **135**) for memory management within the same memory device package. An example of a managed memory device is a managed NAND (MNAND) device.

(40) The memory sub-system **110** includes an address translation circuit **113** and an address translation cache (or ATC **116**) that can be used to perform caching of host memory address translation data used for queues, physical page regions (PRPs) scatter gather lists (SGLs), and data transfer in the memory sub-system **110**. For example, the address translation circuit **113** can receive an address translation request from a HIF circuit that is handling a memory command, request the TA perform a translation of a virtual address located within the memory command, and upon receiving the physical address, store a mapping between the virtual address and the physical address (also referred to as L2P mapping) in the ATC **116**. Upon receiving a subsequent address

translation request that contains the same virtual address, the address translation circuit **113** can verify that the virtual address is a hit at the ATC **116** and directly copy the corresponding address translation from the ATC to a pipeline of the address translation circuit **113** that returns the corresponding address translation to the requesting HIF circuit. Similar caching of address translations can also be performed for DMA operations, which will be discussed in more detail. Further details with regards to the operations of the address translation circuit **113** and the ATC **116** are described below.

(41) FIG. 2 is a schematic block diagram of a system **200** (or device) implementing peripheral component interface express (PCIe) and non-volatile memory express (NVMe) functionality within which the disclosed page request interface support operates in accordance with some embodiments. In various embodiments, the system **200** includes a host system **220** (such as the host system **120**), a memory sub-system **210** (such as the memory sub-system **110**) that in turn includes a controller **215** (such as the controller **115**), or more memory device(s) **130**, and DRAM **222**. In some embodiments, aspects (to include hardware and/or firmware functionality) of the controller **215** is included in the local media controller **135**.

(42) In embodiments, the host system **220** includes a central processing unit (CPU) **209** connected to a host memory **212**, such as DRAM or other main memories. An application program may be stored to memory space **212** for execution by components of the host system **220**. The host system **220** includes a bus **207**, such as a memory device interface, which interacts with a host interface **208**, which may include media access control (MAC) and physical layer (PHY) components, of memory device **130** for ingress of communications from host system **220** to memory device **130** and egress of communications from memory device **130** to host system **220**. Bus **207** and host interface **208** operate under a communication protocol, such as a Peripheral Component Interface Express (PCIe) serial communication protocol or other suitable communication protocols. Other suitable communication protocols include Ethernet, serial attached SCSI (SAS), serial AT attachment (SATA), any protocol related to remote direct memory access (RDMA) such as Infiniband, iWARP, or RDMA over Converged Ethernet (RoCE), and other suitable serial communication protocols.

(43) Memory device **130** may also be connected to host system **220** through a switch or a bridge. A single host system **220** is shown connected with the memory device **130**, and the PCI-SIG Single Root I/O Virtualization and Sharing Specification (SR-IOV) single host virtualization protocol supported as discussed in greater detail below, where the memory device **130** may be shared by multiple hosts, where the multiple hosts may be a physical function **211** (PF) and one or more virtual functions **205** (VFs) of a virtualized single physical host system. In other embodiments, it is contemplated that the SR-IOV standard for virtualizing multiple physical hosts may be implemented with features of the disclosed system and method.

(44) In embodiments, the non-volatile memory arrays (or NVM **206**) of memory device **130** may be configured for long-term storage of information as non-volatile memory space and retain information after power on/off cycles. In the same manner as described with respect to FIG. 1, NVM **206** in FIG. 2 can include one or more dice of NAND type flash memory or other memory discussed with reference to FIG. 1.

(45) The memory sub-system **210** includes a controller **215** (e.g., processing device) which manages operations of memory device **130**, such as writes to and reads from NVM **206**. Controller **215** may include one or more processors **217**, which may be multi-core processors. Processors **217** can handle or interact with the components of memory device **130** generally through firmware code.

(46) Controller **215** may operate under NVM Express (NVMe) protocol, but other protocols are applicable. The NVMe protocol is a communications interface/protocol developed for SSDs to operate over a host and a memory device that are linked over a PCIe interface. The NVMe protocol provides a command queue and completion path for access of data stored in memory device **130** by

host system **220**.

(47) Controller **215** also includes a controller memory buffer (CMB) manager **202**. CMB manager **202** may be connected to the DRAM **222**, to a static random access memory (SRAM) **224**, and to a read-only memory (ROM) **226**. The CMB manager **202** may also communicate with the NVM **206** through a media interface module **228**. The DRAM **222** and SRAM **224** are volatile memories or cache buffer(s) for short-term storage or temporary memory during operation of memory device **130**. In some embodiments, SRAM **224** includes tightly-coupled memory as well. Volatile memories do not retain stored data if powered off. The DRAM generally requires periodic refreshing of stored data while SRAM does not require refreshing. While SRAM typically provides faster access to data than DRAM, it may also be more expensive.

(48) Controller **215** executes computer-readable program code (e.g., software or firmware) executable instructions (herein referred to as “instructions”). The instructions may be executed by various components of controller **215**, such as processor **218**, logic gates, switches, application specific integrated circuits (ASICs), programmable logic controllers, embedded microcontrollers, and other components of controller **215**.

(49) The instructions executable by the controller **215** for carrying out the embodiments described herein are stored in a non-transitory computer-readable storage medium. In certain embodiments, the instructions are stored in a non-transitory computer readable storage medium of memory device **130**, such as in a read-only memory (ROM) or NVM **206**. Instructions stored in the memory device **130** may be executed without added input or directions from the host system **220**. In other embodiments, the instructions are transmitted from the host system **220**. The controller **215** is configured with hardware and instructions to perform the various functions described herein and shown in the figures.

(50) Controller **215** may also include other components, such as a NVMe controller **203**, a media interface module **228** coupled between the NVMe controller **203** and the memory device **130**, and an error correction module **234**. In embodiments, the NVMe controller **203** includes SRAM **204**, an address translation circuit **213** (ATS) having an address translation cache **216**, a direct memory access (DMA) module **230**, a host data path automation (HDPA) circuit **232**, a command parser **236**, a command executor **238**, and a control path **240**. In various embodiments, the address translation circuit **213** is the same as the address translation circuit **113** and the address translation cache **216** is the same as the address translation cache **116**, all of which will be discussed in more detail hereinafter. The SRAM **204** may be internal SRAM of the NVMe controller **203** that is separate from the SRAM **224**. The CMB manager **202** may be directly coupled to the NVMe controller **203** such that the NVMe controller **203** can interact with the CMB manager **202** to access the DRAM **222** and SRAM **224**.

(51) In embodiments, the media interface module **228** interacts with the NVM **206** for read and write operations. DMA module **230** executes data transfers between host system **220** and memory device **130** without involvement from CPU **209**. The HDPA circuit **232** controls the data transfer while activating the control path **240** for fetching PRPs/SGLs, posting completion and interrupts, and activating the DMAs **230** for the actual data transfer between host system **220** and memory device **130**. Error correction module **234** corrects the data fetched from the memory arrays in the NVM **206**. Command parser **236** parses commands to command executor **238** for execution on media interface module **228**.

(52) In embodiments, the controller **215** further includes a page request interface (PRI) handler **219** coupled to or integrated with the processors **217** and the host interface **208**, as will be discussed in more detail. The PRI handler **219** may be processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), firmware (e.g., instructions run or executed on the processors **217**), or a combination thereof.

(53) In various embodiments, the PRI handler **219** automates, from a perspective of the host

interface **208** of the memory controller **215**, sending page miss requests to the host system **120** and detecting and handling page request responses returned from the host system **120**. The page miss requests request a translation agent of the host system **120** to re-pin a physical page of the host memory **212** to a virtual address of each respective page miss request, thus making these pages available again in the host system **220**. In some cases, the re-pinning of the page in the host memory **212** involves finding a valid virtual address to map to an existing physical address (where data is still stored), and completing the mapping for which an address translation can be returned to the controller **215**. Each page request response may include a confirmation that a respective page has been re-pinned in memory as well as the address translation for the pinned page.

(54) FIG. **3** is a schematic diagram illustrating an embodiment of the CMB manager **202** of system **200** of FIG. **2**, but other systems are possible. The CMB manager **202** manages data transactions between host system **220** and a memory device **130** having a controller memory buffer (CMB) **300**. The CMB **300** is a controller memory space which may span across one or more of the DRAM **222**, SRAM **224**, and/or NVM **206**. The contents in CMB **300** typically do not persist across power cycles, so the CMB manager **202** can rebuild the CMB **300** after the system **200** powers on.

(55) One or more types of data structures defined by the NVMe protocol may be stored in the CMB **300** by the CMB manager **202** or may be stored in host memory **212** (FIG. **2**). As described in greater detail below, the host system **220** may initialize the CMB **300** prior to CMB manager **202** storing NVMe data structures thereto. At initialization phase, memory device **130** may advertise to host system **220** the capability and the size of CMB **300** and may advertise which NVMe data structures may be stored into CMB **300**. For example, memory device **130** may store one or more of the NVMe data structures into CMB **300**, including NVMe queues **304** such as submission queues (SQ) **306**, completion queues (CQ) **308**, PRP lists **312**, SGL segments **314**, write data **320**, read data **318**, and combinations thereof.

(56) The NVMe protocol standard is based on a paired submission and completion queue mechanism. Commands are placed by host software into a submission queue (SQ). Completions are placed into the associated completion queue (CQ) by the controller **215**. The host system **220** (or device) may have multiple pairs of submission and completion queues for different types of commands. Responsive to a notification by the host system **220**, the controller **215** fetches the command from the submission queue. Thereafter, the controller **215** processes the command, e.g., performs internal command selection, executes the command (such as performing a write or a read), and the like. After processing the command, the controller **215** places an entry in the completion queue, with the entry indicating that the execution of the command has completed. The controller **215** then generates an interrupt to the host device indicating that an entry has been placed on the completion queue. The host system **220** reviews the entry of the completion queue and then notifies the controller **215** that the entry of the completion queue has been reviewed. As will be discussed in more detail, the address translation circuit **213** may help perform these functions of the controller **215** just discussed.

(57) In general, submission and completion queues are allocated within the host memory **212** where each queue might be physically located contiguously or non-contiguously in the host memory. However, the CMB feature, such as is supported in the NVMe standard, enables the host system **220** to place submission queues, completion queues, physical page region (PRP) lists, scatter gather list (SGL) segments and data buffers in the controller memory rather than in the host memory **212**.

(58) The controller **215** (FIG. **2**) also generates internal mapping tables **221** for use by the CMB manager **202** to map PF and VF data to the correct CMB locations in controller memory DRAM **222** or SRAM **224**. The mapping table **221** itself is typically stored in flip-flops or the DRAM **222** to reduce or eliminate any latency issues. In one implementation, the mapping table may have entries for the PF and each VF, for example.

(59) The NVMe standard supports an NVMe virtualization environment. Virtualized environments

may use an NVM system with multiple controllers to provide virtual or physical hosts (also referred to herein as virtual or physical functions) direct input/output (I/O) access. The NVM system includes of primary controller(s) and secondary controller(s), where the secondary controller(s) depend on primary controller(s) for dynamically assigned resources. A host may issue the Identify command to a primary controller specifying the Secondary Controller List to discover the secondary controllers associated with that primary controller. The SR-IOV defines extensions to PCI Express that allow multiple System Images (SIs), such as virtual machines running on a hypervisor, to share PCI hardware resources (see FIG. 4).

(60) A physical function (PE) is a PCIe function that supports the SR-IOV capability, which in turn allows it to support one or more dependent virtual functions (VFs). These PFs and VFs may support NVMe controllers that share an underlying NVM subsystem with multi-path and namespace sharing capabilities. In such a virtualization environment, the physical function, sometimes referred to as the primary, function, and each virtual function is allocated its own CMB that is a portion of the total controller memory available for CMB use. As used herein, the term physical function refers to a PCIe function that supports SR-IOV capabilities where a single physical host is divided into the physical function and multiple virtual functions that are each in communication with the controller of the memory device. The terms physical function and primary function may be used interchangeably herein.

(61) In an embodiment, the controller **215** advertises the CMB **300** availability only to the physical function (PF) of a virtualized host system such as the host system **220**, where a virtualized host system has a single physical function **211** and one or more virtual functions **205** (or VFs). Also, the advertised CMB **300** availability may be in the form of a total CMB size available for all functions (physical and any virtual functions) such that the physical function **211** may selectively assign itself and all other virtual functions **205** any desired portion of the advertised total CMB size available.

(62) The controller **215** may then store the physical function selected portions of the available CMB **300** in NVMe registers dedicated to each physical function **211** and virtual function, respectively. The virtual function may store a different relative portion size of the advertised CMB size in each NVMe register to account for the different needs the physical function **211** sees for itself and each virtual function. Once the physical function **211** assigns the different amounts and regions of the advertised CMB available for host access (e.g. for direct access by the primary and virtual functions) during the initiation stage, these settings may be managed by the controller **215** to provide access to the respective primary or virtual functions during operations of the memory device.

(63) Controller buffer manager **202** may include a transaction classifier module **322** to classify received host write transactions to CMB **300**. Host write transactions to CMB **300** may be associated with host write command and host read commands. In certain embodiments, transaction classifier module **322** may classify the host write transactions into one of the three NVM data structure groups of NVMe queues **304**, pointers **310**, and data buffers **316**. NVMe queues **304** include host submission queues (SQs) **306** and host completion queues (CQs) **308**. Pointers **310** may include physical region pages (PRP) lists **312** and scatter gather list (SOL) segments **314**. PRP lists **312** contain pointers indicating physical memory pages populated with user data or going to be populated with user data, such as for read or write commands in NVMe queues **304**. SOL segments **314** include pointers indicating the physical addresses of host memory **212** in which data should be transferred from for write commands and in which data should be transferred to for read commands. Data buffers **316** may contain write data **320** to be written to NVM **206** associated with a write command contain and/or read data **318** from memory device **130** associated with a read command.

(64) In certain embodiments, NVMe queues **304**, pointers **310**, and data buffers **316** associated with a particular command may be stored in the CMB **300** by CMB manager **202** to reduce command execution latency by the memory device **130**. For example, a host command entry

written to SQs-implemented CMB **300** avoids fetching the host command entry through the PCIe fabric which may include multiple switches if the SQ is located in the host memory **212**. PRP lists **312** and SGL segments **314** written to CMB **300** of memory device **130** avoids a separate fetch of the PRP lists **312** and SGL, segments **314** through the PCIe fabric if the PRP lists and SGL segments are located in host memory space. Write data **320** written to CMB **300** of memory device **130** avoid having memory device **130** fetch the write data from host memory **212**.

(65) The address translation circuit **213** may communicate through the host interface **208** with the host system **220** and components of the memory sub-system **210**. The address translation circuit **213** may also be incorporated, at least in part, within the host interface **208**, as will be discussed in more detail. The address translation circuit **213** may also retrieve commands from SQs **306**, handle the commands to include retrieving the address translation from the ATC **216**, if present, and submit a completion notification to the CQs **308** for the host system **220**. Thus, in at least some embodiments, the address translation circuit **213** may include or be integrated with the command parser **236**, the command executor **238**, and the DMAs **230**.

(66) FIG. **4** is an example physical host interface between a host system **220** and a memory sub-system implementing caching host memory address translation data in accordance with some embodiments. In at least some embodiments, the example physical host interface also implements NVMe direct virtualization, as was discussed with reference to FIGS. **2-3**. In one embodiment, a controller **415** of the memory sub-system is coupled to host system **220** over a physical host interface, such as a PCIe bus **411A**. In one embodiment, a NVMe control module **405** running on the controller **415** generates and manages a number of virtual NVMe controllers **402-408** within the controller **415**. The virtual NVMe controllers **402-408** are virtual entities that appear as physical controllers to other devices, such as the host system **220**, connected to PCIe bus **411A** by virtue of a physical function **412-418** associated with each virtual NVMe controller **402-408**. FIG. **4** illustrates three virtual NVMe controllers **402-408** and three corresponding physical functions **412-418**. In other embodiments, however, there may be any other number of NVMe controllers, each having a corresponding physical function.

(67) Each of the virtual NVMe controllers **402-408** manages storage access operations for the underlying memory device **130**. For example, virtual NVMe controller **402** may receive data access requests from host system **220** over PCIe bus **411A**, including requests to read, write, or erase data. In response to the request, virtual NVMe controller **402** may identify a physical memory address in memory device **130** pertaining to a virtual memory address in the request, perform the requested memory access operation on the data stored at the physical address and return requested data and/or a confirmation or error message to the host system **220**, as appropriate. Virtual NVMe controllers **404-408** may function in the same or similar fashion with respect to data access requests for one or more memory device(s) **130**.

(68) In embodiments, a NVMe control module **405** associates one of physical functions **412-418** with each of virtual NVMe controllers **402-408** in order to allow each virtual NVMe controller **402-408** to appear as a physical controller on the PCIe bus **411A**. For example, physical function **412** may correspond to virtual NVMe controller **402**, physical function **414** may correspond to virtual NVMe controller **404**, and physical function **418** may correspond to virtual NVMe controller **408**. Physical functions **412-418** are fully featured PCIe functions that can be discovered, managed, and manipulated like any other PCIe device, and thus can be used to configure and control a PCIe device (e.g., virtual NVMe controllers **402-408**). Each physical function **412-418** can have some number of virtual functions (VFs) associated therewith. The VFs are lightweight PCIe functions that share one or more resources with the physical function and with virtual functions that are associated with that physical function. Each virtual function has a PCI memory space, which is used to map its register set. The virtual function device drivers operate on the register set to enable its functionality and the virtual function appears as an actual PCIe device, accessible by host system **220** over the PCIe bus **411A**.

(69) In at least some embodiments, the controller **415** is further configured to control execution of memory operations associated with memory commands from the host system **220** at one or more memory devices(s) **130** and one or more network interface cards (NIC(s)) **450**, which are actual physical memory devices. In these embodiments, the controller **415** communicates with the memory devices **130** over a second PCIe bus **411B** and communicates with the NICs **450** over a third PCIe bus **411C**. Each memory device **130** can support one or more physical functions **432** and each NIC **450** can support one or more physical functions **436**. Each physical function **432-436** can also have some number of virtual functions (VFs) associated therewith.

(70) In these embodiments, each physical function **412-418** and **432-436** can be assigned to any one of virtual machines VM(**0**)-VM(**n**) in the host system **220**. When I/O data is received at a virtual NVMe controller **402-408** from a virtual machine, a virtual machine driver (e.g., NVMe driver) provides a guest physical address for a corresponding read/write command. The NVMe control module **405** can translate the physical function number to a bus, device, and function (BDF) number and then add the command to a direct memory access (DMA) operation to perform the DMA operation on the guest physical address. In one embodiment, the controller **115** further transforms the guest physical address to a system physical address for the memory sub-system **110**.

(71) Furthermore, each physical function **412-418** and **432-436** can be implemented in either a privileged mode or normal mode. When implemented in the privileged mode, the physical function has a single point of management that can control resource manipulation and storage provisioning for other functions implemented in the normal mode. In addition, a physical function in the privileged mode can perform management options, including for example, enabling/disabling of multiple physical functions, storage and quality of service (QoS) provisioning, firmware and controller updates, vendor unique statistics and events, diagnostics, secure erase/encryption, among others. Typically, a first physical function can implement a privileged mode and the remainder of the physical functions can implement a normal mode. In other embodiments, however, any of the physical functions can be configured to operate in the privileged mode. Accordingly, there can be one or more functions that run in the privileged mode.

(72) The host system **220** can run multiple virtual machines VM(**0**)-VM(**n**), by executing a software layer, often referred to as a hypervisor **424**, above the hardware and below the virtual machines, as schematically shown in FIG. **4**. In one illustrative example, the hypervisor **424** may be a component of a host operating system **422** executed by the host system **220**. Alternatively, the hypervisor **424** may be provided by an application running under the host operating system **422**, or may run directly on the host system **220** without an operating system beneath the hypervisor **424**. The hypervisor **424** may abstract the physical layer, including processors, memory, and I/O devices, and present this abstraction to virtual machines VM(**0**)-VM(**n**) as virtual devices, including virtual processors, virtual memory, and virtual I/O devices. Virtual machines VM(**0**)-VM(**n**) may each execute a guest operating system which may utilize the underlying virtual devices, which may, for example, map to the memory device **130** or the NIC **450** managed by one of virtual NVMe controllers **402-408** in the memory sub-system. One or more applications may be running on each VM under the guest operating system.

(73) In various embodiments, each virtual machine VM(**0**)-VM(**n**) may include one or more virtual processors and/or drivers. Processor virtualization may be implemented by the hypervisor **424** scheduling time slots on one or more physical processors such that from the perspective of the guest operating system, those time slots are scheduled on a virtual processor. Memory virtualization may be implemented by a page table (PT) which is a memory structure translating guest memory addresses to physical memory addresses. The hypervisor **424** may run at a higher privilege level than the guest operating systems, and the latter may run at a higher privilege level than the guest applications.

(74) In one embodiment, there may be multiple partitions on host system **220** representing virtual machines VM(**0**)-VM(**n**). A parent partition corresponding to virtual machine VM(**0**) is the root

partition (i.e., root ring **0**) that has additional privileges to control the life cycle of other child partitions (i.e., conventional ring **0**), corresponding, for example, to virtual machines VM(**1**) and VM(**n**). Each partition has corresponding virtual memory, and instead of presenting a virtual device, the child partitions see a physical device being assigned to them. When the host system **220** initially boots up, the parent partition can see all of the physical devices directly. The pass through mechanism (e.g., PCIe Pass-Through or Direct Device Assignment) allows the parent partition to assign an NVMe device (e.g., one of virtual NVMe controllers **402-408**) to the child partitions. The associated virtual NVMe controllers **402-408** may appear as a virtual storage resource to each of virtual machines VM(**0**), VM(**1**), VM(**n**), which the guest operating system or guest applications running therein can access. In one embodiment, for example, virtual machine VM(**0**) is associated with virtual NVMe controller **402**, virtual machine VM(**1**) is associated with virtual NVMe controller **404**, and virtual machine VM(**n**) is associated with virtual NVMe controller **408**. In other embodiments, one virtual machine may be associated with two or more virtual NVMe controllers. The virtual machines VM(**0**)-VM(**n**), can identify the associated virtual NVMe controllers using a corresponding bus, device, and function (BDF) number, as will be described in more detail below. (75) In some embodiments, the hypervisor **424** also includes a storage emulator **426** coupled to the NVMe drivers on the virtual machines VM(**0**)-VM(**n**) and that is coupled with a physical function NVMe driver **432** of the hypervisor **424**. The physical function NVMe driver **432** can drive, with the help of the virtual NVMe controllers **402-408**, the physical functions **412-418** over the PCIe bus **411A** and also drive the physical functions **432** available on the memory devices **130** over the PCIe bus **411A** and the second PCIe bus **411B**. Further, the hypervisor can include a NIC emulator **428** coupled to NIC drivers on the virtual machines VM(**0**)-VM(**n**) and that is coupled with a physical function NIC driver **436** of the hypervisor **424**. The physical function NIC driver **434** controls the PFs **436** of the NICs **450** over the PCIe bus **411A** and the third PCIe bus **411C** in various embodiments.

(76) In at least some embodiments, the host system **220** submits memory commands (e.g., erase (or unmap), write, read) to a set of submission queues **442**, input/output (I/O) commands to a set of I/O queues **446**, and administrative (“admin”) commands to a set of admin queues **448**, which are stored in the host memory **212** of the host system **220** or in one of the CMB **300** (FIG. 3). The controller **415** can retrieve these memory commands over the PCIe bus **411A** and handle each memory command in turn, typically according to a priority (such as handling reads in front of writes). When the controller **415** has completed handling a memory command that resides in the set of submission queues **442**, the controller **415** returns an acknowledgement of memory command completion by submitting a completion entry in a corresponding completion queue of a set of completion queues **444**, which are also stored in the host memory **212**. In some embodiments, the host memory **212** is composed of DRAM or other main memory type memory. In various embodiments, the queues **442**, **444**, **446**, **448** can number into the hundreds (or thousands) and are ordered sequentially (e.g., contiguously) according to virtual addresses of the memory commands. In other words, the queuing of memory commands within these queues is ordered sequentially based on the virtual addresses in those memory commands within a virtual memory space.

(77) In disclosed embodiments, the controller **415** further includes an address translation circuit **413**, which includes an ATC **416** (or “cache”) similarly introduced with reference to FIGS. 1-3. The ATC **416** may be static random access memory (SRAM), tightly-coupled memory (TCM), or other fast-access memory appropriate for use as cache. The address translation circuit **413** can be coupled to NVMe control module **405** and generally coupled to PCIe protocol components (such as the virtual NVMe controllers) of the controller **415**. In this way, the address translation circuit **413** provides host interface circuitry that facilitates obtaining address translations and other handling of memory commands retrieved from the queues **442**, **444**, **448**. The address translation services (ATS) of the address translation circuit **413** further enables direct connection between the one or more memory devices **130** and the virtual machines VM(**0**)-VM(**n**) to achieve a near PCIe line rate

of data communication performance, which bypasses the hypervisor **424**.

(78) As explained, the address translation circuit **413** can receive (or retrieve) an address translation request from a HIF circuit that is handling a memory command, request the TA perform a translation of a virtual address located within the memory command, and upon receiving the physical address, store a mapping between the virtual address and the physical address (also referred to as L2P mapping) in the ATC **416**. Upon receiving a subsequent address translation request that contains the same virtual address, the address translation circuit **413** can verify that the virtual address is a hit at the ATC **416** and directly copy the corresponding address translation from the ATC to a pipeline of the address translation circuit **413** that returns the corresponding address translation to the requesting HIF. Similar caching of address translations can also be performed for DMA operations, which will be discussed in more detail. Accordingly, the functionality of the address translation circuit **413** can enable, for many address translation requests, bypassing any need to interact with the hypervisor **424**, which being software, is the bottleneck and slows performance of obtaining address translations in the absence of caching such address translations. The resultant speed, latency, and throughput performance increases through the ATS functionality can be significant.

(79) FIG. 5 is a system **500** in which the memory sub-system controller, such as a PCIe controller **515**, contains an address translation cache (ATC) **516** in accordance with some embodiments. Within the system **500**, a host system **520** (or any host system discussed herein) includes a translation agent (TA) **517** and an address translation and protection table (ATPT) **519**, together which can be employed by the host system **520** to provide address translations to the PCIe controller **515** according to PCIe protocol. Specifically, the TA **517** can return a physical page address in response to a submitted virtual address in an address translation request or a “no translation” in the case the corresponding physical page has been swapped out and thus there is not current translation for the virtual address.

(80) These address translations can be associated with memory commands resident in the host memory **512** (and/or CMB **300**) being handled by the PCIe controller **515** (or by other PCIe device or virtual PCIe device). To provide translations, the TA **517** is configured to communicate with address translation requests/responses through a PCIe root complex **523**. In some systems, the TA **517** is also known as an input/output memory management unit (IOMMU) that is executed by the hypervisor **424** or virtual machine manager running on the host system **520**. Thus, the TA **517** can be a hardware component or software (IOMMU) with a dedicated driver.

(81) In various embodiments, to avoid such increased I/O traffic over a PCIe bus **511A** between the host system **520** and the PCIe controller **515**, the address translations provided to the PCIe controller **515** can be cached in the ATC **516** and accessed to fulfill later (or subsequent) ATS-generated address translation requests without having to go back to the TA **517** with renewed requests for each needed translation. If entries in the ATC **516** are invalidated due to assignment changes between the virtual and physical addresses within the TA **517** and the ATPT **519**, then the PCIe controller **515** (e.g., the ATS in the PCIe controller **515**) can purge corresponding entries within the ATC **516** and in other ATS queues in host interface circuitry.

(82) FIG. 6 is a memory sub-system **610** for page request interface support in caching host memory address translation data for multiple host interface circuits in accordance with some embodiments. In various embodiments, the memory sub-system **610** can be the same as the memory sub-system **110** of FIG. 1 or the memory sub-system **210** of FIG. 2 or that of FIG. 5. In these embodiments, the memory sub-system **610** includes a controller **615** having a controller memory buffer (CMB) **630** and the PRI handler **219** (discussed with reference to FIG. 2). In some embodiments, the controller **615** is the controller **115**, **215**, or **415**, the local media controller **135**, or a combination thereof (see FIG. 1). In at least one embodiment, the CMB **630** is the CMB **300** discussed with reference to FIG. 3, and thus may also include NVM of the memory device **130**.

(83) In various embodiments, the CMB **630** includes host Advanced eXtensible Interface (AXI)

interfaces **631**, CMB control/status register(s) **632**, host read command SRAM/FIFO **634**, host data SRAM **636**, CMB host write buffer **638**, controller write buffer **640**, host read buffer **644**, controller data SRAM **646**, controller read buffer **648**, and a controller read command SRAM/FIFO **650** which support the functionality of the controller **615**, wherein FIFO stands for first-in-first-out buffer.

(84) In some embodiments, the controller **615** includes a PCIe system-on-a-chip (SoC) **601** and host interface circuitry **603**. The PCIe SoC **601** may include PCIe IP that facilitates communication by the host interface circuitry **603** with the host system **620** (including a TA **617**) using PCIe protocols. Thus, the PCIe SoC **601** may include capability and control registers present for each physical function and each virtual function of the memory sub-system **610**. The PCIe SoC **601** may include integrated development environment (IDE) link encryption circuit **605**, which encrypts translation layer packets (TLPs) passed over a PCIe bus **611A** to a host system **620** and to decrypt TLPs received from the host system **620**.

(85) In various embodiments, the host interface circuitry **612** includes the local memory **619** (SRAM, TCM, and the like), a number of host interface circuits **608A**, **608B**, **608C**, . . . **608N** (or together host interface circuits **608**), an address translation circuit **613** (or ATS) that includes an ATC **616** and translation logic **602**, and an AXI matrix **623** employable to send interrupts to a host system **620** in interaction with reference to handling memory commands and DMAs. The host interface circuits **608A-608N**, also referred to herein as HIF circuits **608**, include hardware components that help to fetch and process commands (for host queues) and data (for DMAs) to perform a particular function in command and DMA processing, individual ones of the host interface circuits **608A-608N** at times request an address translation of a virtual address or a guest physical address (e.g., associated with a command or DMA). In embodiments, to do so, individual HIF circuits **608** request the address translation circuit **613** to provide the translation. In this way, the address translation circuit **613** interfaces with and supports both address translation generation and invalidation on behalf of the respective host interface circuits **608A-608N**, as will be discussed in more detail hereinafter.

(86) Further, in at least one embodiment, the host interface circuit **608N** may be a processing bridge circuit **608N** that provides a hardware interface between the PM handler **219** and the other host interface circuits **608A**, **608B**, **608C**, etc. The processing bridge circuit **608N** may include various buffers and logic for processing page miss requests, page request responses, and other commands and messages that will be discussed in more detail with reference to FIG. **10**.

(87) The address translation circuit **613** can employ many fields and parameters, such as a smallest translation unit (STU) of data (defined by a particular size typically smaller than a host tag) and an invalidation queue depth, beyond which queue entries are purged. The translation logic **602** may use PCIe memory read TLPs to request a translation for a given untranslated address from the TA **617** (or IOMMU). These address translation requests may carry separate PCIe tags and identifiers, for example. The PCIe tags may include, for example “PCIe Memory read completion,” indicating the translated address and the size of translation, “S” bit informing the size of translation (e.g., a translated address may represent contiguous physical address space of 128 KB size or similar size), “R” and “W” bits provide Read/write permissions to a page in physical address space, and “U” bit tells the device to do DMA using the untranslated address only. The U bit may be helpful when buffers are one-time use and TA **617** does not have to send invalidations to the controller **615**.

(88) In embodiments, when a translation changes in the TA **617**, the ATC **616** (one or more caches) in the memory sub-system **610** should purge the old entries corresponding to a virtual address that has been invalidated. Invalidation requests may be sent by the TA **617** to the address translation circuit **613** using PCIe message TLPs, which requests may be directed at an HIF circuit that is handling some function of memory command processing. The address translation circuit **613** can direct the HIF circuits to purge all DMA for an address range that is being invalidated, remove such entries from the ATC **616**, and upon confirmation from the HIF circuits of invalidation, send an

invalidation completion (e.g., another PCIe Message TLP) to the host system **620**.

(89) In some embodiments, the address translation circuit **613** stores, in the ATC **616**, an address translation that returns from the TA **617** in response to an address translation request. In various embodiments, this address translation may include an I/O submission queue (SQ) base address, PRP/SGLs of outstanding commands in one or more HIF circuits, and/or I/O completion queue (CQ) base address. The ATC **616** can be configured to handle finding an empty slot in the ATC **616** and storing a new translation in the empty slot, looking up when data TLPs show up, and purging entries during function level reset (FLR), invalidations, and other resets. The ATC **616** can further be configured to age-out older entries to make space when the cache is running full.

(90) In various embodiments, a queue portion **416A** of the ATC **616** cache can be sized to include at least one or two entries per queue of the SQs **442**, CQs **444**, I/O queues **446**, and admin queues **448**, although more are envisioned as cache memory device sizes and costs decrease. With sufficient space for two entries, the ATC **616** can store the address translation of a current page associated with a queue as well as a next page (e.g., that sequentially follows the virtual address of the current page). This look-ahead buffering in the ATC **616** of address translations for a predetermined number of queues enables greatly reducing the number of misses at the ATC while keeping a size of the ATC reasonable given the expense of cache memory. Further, a DMA portion **416B** of the ATC **616** (for storing data for DMAs) could be expanded for some integral multiple (e.g., 2-4 times) the size of the queue portion of the ATC **616**.

(91) In at least some embodiments, the translation logic assigns a queue identifier to each queue of the SQs **442**, CQs **444**, I/O queues **446**, and admin queues **448** (“the queues”). This could be as simple as sequentially incrementing a count number for each subsequent sequentially-ordered queue of a set of queues. The translation logic **602** may then index the queue portion **416A** of the ATC **616** cache according to the queue identifier of the respective address translation stored therein and internally (within the address translation circuit **613**) track address translation requests and responses using such queue identifiers. Further, the translation logic **602** may index the DMA portion **416B** of the cache according to host tag value and internally track DMA-related address translations requests and responses using host tag values.

(92) In these embodiments, each DMA command has a host tag (“htag”) with a virtual address. The translation logic **602** may be configured to store in the ATC **616** one translation per htag for data if the data is more than or equal to a pre-configured size of data. The translation logic **602** may further be configured to store in the ATC **616** one translation per htag for metadata (with no size limit dictating whether to cache). During further operation of the controller **615**, the address translation is used as long as subsequent TUs (within the htag) use the same translated physical address range. The cached translation gets replaced when the DMA command starts transferring to or from a different memory range. The cached address translation also gets replaced when the htag gets assigned to a different command, e.g., that uses a different memory location.

(93) FIG. 7 is an example memory sub-system controller **715** including an address translation circuit **713** (or ATS) implementing caching host memory address translation data in accordance with some embodiments. The address translation circuit **713** may be a more detailed instantiation of the address translation circuit **113**, **213**, **413**, and/or **613** discussed with reference to FIG. 1, FIG. 2, FIG. 4, and FIG. 6, respectively.

(94) In various embodiments, the address translation circuit **713** includes a pipeline of queues, buffers, and multiplexers that facilitate the flow of address translation requests and responses and that interfaces with an address translation cache (ATC) **716**, similarly as introduced and discussed previously. The multiplexers are generally included within the translation logic **602** and are therefore not always individually numbered or discussed separately from the translation logic **602**.

(95) In embodiments, this pipeline begins with address translation requests flowing into a set of request staging queues **702** from host interface (HIF) circuits **608** (see FIG. 6). These address translation requests are often received as a series of address translation requests from any given

HIF circuit **608**, thus each queue in the set of request staging queues **702** includes multiple available entries for multiple virtual addresses associated with an incoming series of address translation requests, which can be received as a group for example from a given HIF circuit **608**. The set of request staging queues **702** buffer the address translation requests, each including a virtual address, which are received from a host interface circuit. A multiplexer **703** then pushes each address translation request into both a set of reordering buffers **714** and to the translation logic **602** that was initially discussed with reference to FIG. **6**. Within the set of reordering buffers **714**, the address translation requests will provide initial entries that will later be completed with address translations, which will be provided in response to corresponding requesting HIF circuits **608**. In embodiments, the translation logic **602** is configured to store incoming address translations, e.g., from the set of reordering buffers **714**, into the ATC **716**.

(96) In embodiments, the translation logic **602** is coupled to the set of request staging queues **702**, the set of reordering buffers **714**, and the ATC **716**, among other components. The translation logic **602** can determine a queue identifier for respective queues of multiple submission queues (SQs) and multiple completion queues (CQs) discussed previously and identify the queue identifier for respective address translation requests. The translation logic **602** can further index the queue portion **416A** of the cache according to the queue identifier of the respective address translations requests.

(97) In at least some embodiments, the translation logic **602** is configured to, for each address translation request in the set of request staging queues, store, in the ATC **716**, a first address translation corresponding to a current page associated with the address translation request and store, in the ATC **716**, a second address translation corresponding to a subsequent page that sequentially follows the current page according to virtual address numbering. This look-ahead buffering in the ATC **716** of address translations for a predetermined number of queues enables greatly reducing the number of misses at the ATC **716** while keeping a size of the ATC reasonable given the expense of cache memory. As cost of cache memory decreases, it is envisioned that one or more additional address translations that sequentially follow the first address translation (according to virtual addresses) can also be stored in the ATC **716**.

(98) In some embodiments, for incoming DMA requests, the translation logic **602** can identify a first host tag, within at least some of the address translation requests, associated with data for DMA. The translation logic can further identify a second host tag, within the at least some of the address translation requests, associated with metadata of the data. The translation logic can then index the DMA portion **416B** of the cache according to host tag value, to include values corresponding to each first host tag and each second host tag. In embodiments, in response to a translation unit (TU) of a subsequent address translation request being within a translated memory range as a cached host tag, the translation logic uses an address translation of the cached host tag in the ATC **716** to satisfy the subsequent address translation request. Further, in response to the subsequent address translation request targeting a different memory range than the translated memory range, the translation logic can evict the address translation from the ATC **716**. In response to the cached host tag being assigned to a different command, the translation logic **602** can also evict the address translation from the cache (ATC **716**).

(99) In some embodiments, the translation logic **602** further, for each incoming address translation request, determines whether the virtual address (or queue identifier) within the address translation request hits or misses at the cache (ATC **716**). If the virtual address has a hit within the cache, the translation logic **602** may reinsert a corresponding address translation (including the mapped physical address) into the reordering buffers **714** to be provided out to the requesting HIF circuit **608**. Thus, address translations that hit at the ATC **716** can be provided back to the requesting HIF circuit **608** at close to line rate without any further delay in requesting the TA/IOMMU at the host system **620** to perform the translation.

(100) In embodiments, the address translation circuit **713** further includes a set of outbound request

queues **704** to buffer one or more of the address translation requests that miss at the cache (the ATC **716**), e.g., as an extra pipeline stage for staging these address translation requests to be sent to the host system **620**. As these address translation requests are forwarded on to the host system **620**, the translation logic **602** pushes each address translation request into a set of pending response queues **708**. In embodiments, the set of pending response queues **708** is configured to buffer respective address translation requests that are waiting for an address translation from the host system **620** while maintaining an order as received within the set of request staging queues **702**. Maintaining this order within the set of pending response queues **708** helps the set of reordering buffers **714** to properly reorder address translations that come back from the host system **620** (e.g., the TA **617** or IOMMU) despite being pushed out of the set of request staging queues **702** while waiting on the TA **617**.

(101) In these embodiments, the address translation requests that missed at the ATC **716** are forwarded to an AXI converter **722**, which sends out address translation requests to the host interface **208** over a set of address channels for reads **732** and destined to the TA **617**. In this way, the address translation circuit **713** can request a translation agent (such as the TA **617**) of the host system to provide physical addresses that map to a first virtual address of the first address translation (for the current page) and to a second virtual address of the second address translation (for the subsequent page that sequentially follows the current page according to virtual address numbering). In response to the TA **617** providing the address translations, the address translations are received over a set of data channels **734** for reads and through an AXI converter **724**.

(102) In disclosed embodiments, the translation logic **602** pops the incoming address translations into the set of pending responses queues **708** while also storing the address translations (e.g., the first address translation and the second address translation) in the set of reordering buffers **714**. The set of reordering buffers **714** have also received, from the set of pending response queues **708**, the proper order of the address translations in relation to the order of receipt of the address translation requests in the set of request staging queues **702**. As mentioned, this may include a series of address translation requests from the same HIF circuit **608**, and thus the set of reordering buffers **714** can be configured to reorder each pair of address translations for each of these address translation requests to match the order the address translation requests were buffered into the set of pending response queues **708**. Thus, the set of reordering buffers **714** can reorder the address translations in this way and provide the first address translation (for each address translation request) to the requesting HIF circuit **608** in the same order as was received with the exception of those that hit at the ATC **716** and have already been supplied to the requesting HIF circuit **608**. In these embodiments, the translation logic **602** detects the incoming address translations and copies them from the set of reordering buffers into the ATC **716**, which can be accessed later for matches with subsequent address translation requests.

(103) In some embodiments, a translation completion response may come back as a “miss” from the host system **620** because the page corresponding to the virtual address has been unmapped. The translation logic **602** can then reply to the requesting HIF circuit **608** that the virtual address of a particular address translation requested missed. In response to receiving a TLP of such a miss, the HIF circuit **608** may be triggered to generate a page request interface (PRI) message to the host system **620** to remap the virtual address to a physical page having a physical address.

(104) In at least some embodiments, the address translation circuit **713** further includes a set of invalidation queues **718** to buffer invalidation requests received from a translation agent of the host system. More specifically, the invalidation requests may be received over a set of data channels for writes **736** into an AXI converter **726** that receives the invalidation requests from the host system **620**. The translation logic **602** can then place each incoming invalidation request into the set of invalidation queues **718** to be handled in an order received.

(105) In these embodiments, the address translation circuit **713** further includes invalidation handler logic **750** coupled to the set of invalidation queues **718**, to the host interface circuits **608**, to

the set of request staging queues **702**, to the set of pending response queues **708**, and to the set of reordering buffers **714**. In embodiments, the invalidation handler logic **750** detects an invalidation request within the set of invalidation queue **718**, the invalidation request corresponding to a virtual address of the virtual addresses. The invalidation handler logic **750** can further cause address translations associated with the virtual address to be marked as invalid within the set of request staging queues **702**, the set of pending response queues **704**, the set of reordering buffers **714**, and the cache (the ATC **716**). In embodiments, the set of outbound request queues **704** triggers (e.g., via the translation logic **602**) the set of reordering buffers **714** to mark an entry therein as invalid in response to an invalidation signal received from the invalidation handler logic **750**. The invalidation handler logic **750** can further send associated invalidation requests to the respective host interface circuits **608**, which also would be expected to purge invalid entries.

(106) In some embodiments, the address translation circuit **713** (e.g., the corresponding queues and buffers) remove address translations marked as invalid within the set of request staging queues **702**, the set of pending response queues **704**, the set of reordering buffers **714**, and the cache (ATC **716**). Each HIF circuit **608** can send an invalidation complete message to the address translation circuit **713**, to which the invalidation handler logic **750** and reply with an acknowledgement (ACK). In this way, invalidation handler logic **750** can confirm removal of the address translations by the host interface circuits **608** followed by sending an invalidation completion response to the translation agent (TA **617**). The ACK responses may be used by individual HIF circuits **608** to unfreeze any frozen function arbitration.

(107) More specifically, when the invalidation handler logic **750** has confirmed the address translation circuit **713** and the HIF circuits **608** have purged queue, buffer, and cache entries of invalid address translations associated with an invalidation requests, the invalidation handler logic **750** may send an invalidation response to an AXI converter **728**. The AXI converter **728** can place the invalidation response out on a set of data channels for writes **738** to the TA **617** at the host system **620**.

(108) FIG. **8** illustrates a flow chart of an example method **800** of caching host memory address translation data in a memory sub-system in accordance with some embodiments of the present disclosure. The method **800** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **800** is performed by the address translation circuit **113**, **213**, **413**, **613**, **713** of FIG. **1**, FIG. **2**, FIG. **4**, FIG. **6**, and FIG. **7**, respectively. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

(109) At operation **810**, the processing logic buffers, within a set of request staging queues of an address translation circuit, address translation requests received from host interface circuits and each comprising a virtual address.

(110) At operation **820**, the processing logic stores, within a set of pending response queues, respective address translation requests that are waiting for an address translation from a host system while maintaining an order as received within the set of request staging queues.

(111) At operation **830**, the processing logic reorders, within a set of reordering buffers, address translations according to the order maintained within the set of pending response queues, wherein each address translation comprises a physical address mapped to a respective virtual address.

(112) At operation **840**, the processing logic sends, from the reordering buffers, the address translations to a corresponding host interface circuit of the host interface circuits that sent a

corresponding address translation request.

(113) At operation **850**, the processing logic stores, in a cache coupled with the set of request staging queues and the set of reordering buffers, a plurality of the address translations, associated with the address translations requests, for future access by the host interface circuits. Also, at operation **850**, the processing logic can further store a plurality of pointers (e.g., SGS/SGL pointers) for outstanding direct memory address (DMA) command within on or more host interface circuits.

(114) At operation **860**, when applicable, the processing logic reinserts, into the set of reordering buffers, a first address translation from the cache for a subsequent request for the first address translation by a host interface circuit.

(115) FIG. **9** is a block diagram of a system **900** for host interface circuits interacting with the PRI handler **219** in accordance with some embodiments. In some embodiments, the system **900** includes the address translation circuit **613** (see FIG. **6**), a host interface circuit **608** (such as any of host interface circuits **608A**, **608B**, **608C**, . . . **608N** discussed with reference to FIG. **6**), the processing bridge circuit **608N** (see FIG. **10**), the PRI handler **219**, and a controller memory **912**.

(116) In various embodiments, a host memory manager (such as the TA **617**) can swap out a page of the host memory **212** for a variety of reasons, which can be done without informing user processes or any of the virtual machines VM(**0**)-VM(**n**). As was discussed, a device, including the controller **615** that employs the address translation circuit **613**, cannot assume that the page that is needed to DMA is present. Accordingly, a page request interface, which will be discussed in more detail with reference to FIG. **10**, in conjunction with the PRI handler **219** can provide the ability for the host interface circuitry **603** to cause the translation agent **617** (e.g., the host memory manager) to re-pin a page of memory, return a confirmation that the page has been re-pinned, and provide the corresponding address translation.

(117) In these embodiments, the host interface circuit **608** includes a translation request queue **904**, a translation response queue **908**, a hardware pipeline **907** having a number of hardware stages (HW_**0** . . . HW_**N**), a command packet queue **905** that buffers incoming command packets to be processed by the hardware pipeline **907**, and a set of control registers **911**. The set of control registers **911** may be present in relation to a physical function (PH) where related or other virtual functions (VFs) may use this set of controller registers **911** as well. Values within the set of controller registers **911** may define an allocation for each PF, how many PRI-related requests can be outstanding for a given PF, and other settings or parameters for hardware functionality of the PF.

(118) In various embodiments, while a command packet is entering or at one of the hardware stages of the hardware pipeline **907**, the host interface circuit **608** submits a virtual address of the command packet to the translation request queue **904** to request that the address translation circuit **613** provide an address translation required for the hardware pipeline **907** to complete processing the command packet. If the address translation circuit **613** returns an address translation that was previously stored in the ATC **616**, e.g., via the translation response queue **908**, the host interface circuit **608** proceeds normally with no pause in processing the command packet.

(119) In at least some embodiments, the host interface circuit **608** receives a message from the address translation circuit **613**, e.g., via the translation response queue **908**, that the address translation request has missed at the cache or ATC **616**, and thus is a host interface circuit **608** affected by the miss at the cache. In embodiments, this translation miss message triggers the host interface circuit **608**, e.g., by referencing values within the set of control registers **911**, to remove the command packet from the hardware pipeline **907**. The command packet can be associated with a function that relies on the address translation corresponding to the address translation request that missed at the cache. The host interface circuit can further send a translation miss message to the PM handler **219** (e.g., via the controller memory **912**), the translation miss message including the virtual address that was in the original address translation request. This translation miss message may be queued within controller memory **912** before arriving at the PM handler **219**, as will be

discussed with reference to FIG. 10.

(120) In various embodiments, the translation miss message includes information that identifies where and/or how to reintroduce the command packet into the hardware pipeline. For example, in hardware stage HW_2 triggered the address translation request, and the command packet was removed at this third hardware stage of the hardware pipeline 907, the information included in the translation miss message to the PM handler 219 can be an identifier of this third hardware stage (HW_2). In embodiments, the set of controller registers 911 stores hardware stage identifiers that are mapped to the respective hardware stages of the hardware pipeline 907. In response to a restart message received for the PRI handler 219 corresponding to the translation miss message that was previously sent to the PRI handler 219, the host interface circuit 608 can reintroduce a command packet into the hardware pipeline 907 where the command packet had been originally removed. In this way, DMAs can proceed (even if slightly delayed during PRI-related processing) and the host interface circuitry is made to more efficiently continue to process commands and other actions in interacting with the host memory 212.

(121) FIG. 10 is a schematic block diagram of a system 1000 (e.g., a processing bridge system) including the processing bridge circuit 608N implementing page request interface support in caching host memory address translation data in a memory sub-system in accordance with some embodiments. In these embodiments, the system 1000 include the processing bridge circuit 608N, a number of queues stored in a controller memory 1012 (e.g., of the controller 615), the PRI handler 219, and an AXI matrix 1023. In some embodiments, the controller memory 1012 is DRAM, but can also be SRAM or combination of volatile memory. The controller memory 1012 may be the same as the controller memory 912 (FIG. 9). In these embodiments, the processing bridge circuit 608N is coupled between the controller memory 1012 and the AXI matrix 1023, which provides data and address communication to the host system 620, and also has access to the address translation circuit 613 previously discussed herein. Reference to the host system 620 may be understood to include reference to any other host system herein, including the host system 220 of FIG. 2.

(122) In various embodiments, the PRI handler 219 is coupled to the controller memory 1012 and is optionally integrated with the processing bridge circuit 608N. As discussed, the PRI handler 219 may be processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), firmware (e.g., instructions run or executed on the processors 217), or a combination thereof. The controller memory 1012 can store a number of different queues to facilitate communication between the PRI handler 219 and the processing bridge circuit 608N.

(123) In embodiments, for example, the controller memory 1012 stores a set of translation miss queues 1070 coupled to respective HIF circuits 608. For example, the host interface circuit 608A can be coupled to a HIF1 circuit translation miss queue 1070A, the host interface circuit 608B can be coupled to a HIF2 circuit translation miss queue 1070B, the host interface circuit 608C can be coupled to a HIF3 circuit translation miss queue 1070C, and a further host interface circuit can be coupled to a HIF4 circuit translation miss queue 1070D, where additional or fewer miss request queues are envisioned. Each queue of the set of translation miss queues 1070 receives and buffers page miss messages from respective HIF circuits 608, where each page miss message is indicative of a miss at the address translation cache (ATC) 616 for a host address translation. The controller memory 1012 can further store an outbound queue 1072, a host request queue 1074, and a bridge translation miss queue 1075, which will be discussed in more detail.

(124) In various embodiments, the PRI handler 219 tracks activities associated with each page miss message received at the set of translation miss queues 1070 from respective ones of the host interface circuits 608. In these embodiments, the PRI handler 219 tracks translation miss messages received from the host interface circuits 608, each translation miss message including a virtual address of a miss at the cache (or ATC 616). The PRI handler 219 further removes duplicate

translation miss messages having an identical virtual address, e.g., so as not to send multiple duplicate page miss requests to the host system **620** unnecessarily. The PRI handler **219** further creates a set of page miss requests from non-duplicate translation miss messages that are categorized into page request groups. Thus, the page miss requests are associated with page miss messages received from a respective host interface circuit. The PRI handler **219** further queues the page request groups to be sent to a translation agent of the host system, e.g., within the host request queue **1074**.

(125) In these embodiments, each page request group, for example, corresponds to a host interface circuit of the host interface circuits **608**. Each page miss request causes the translation agent (e.g., TA **617**) to re-pin a physical page of memory to the virtual address of a respective page miss request. In at least some embodiments, the page miss requests are PCIe message TLPs that contain a virtual address of the page and optionally a process address space identifier (PASID). The PASID may be included when the host system **620** is executing processes (such as the virtual machines VM(0)-VM(n)) that share the controller **615** or memory devices controlled by the controller **615**.

(126) In some embodiments, each page needs its own page miss request, and potentially, there can be several thousand page miss requests for a big I/O load, such as may be required for interacting with a NIC. Each page miss request may invoke page handler software at the host system **620** that does batch processing of multiple page requests. For simplicity of illustration, the TA **617** (which can be an IOMMU) may be understood to include such page handler software. In embodiments, sending the same page address in multiple page miss requests may not trigger an error, although doing so may burn host CPU cycles.

(127) In these embodiments, the PRI handler **219** further categorizes the page miss requests into the page request groups, e.g., according to an incoming queue of the set of translation miss queues **1070** or another identifier such as a hardware identifier of a source of the page miss requests. The PRI handler **219** can further restrict the page miss requests that are outstanding for a particular host interface circuit **608** to a threshold number of page miss requests. This may be permed in firmware or as practical limitation based on an allocation of the controller memory **1012** to a respective translation miss queue assigned to the particular host interface circuit **608**.

(128) In various embodiments, the outbound queue **1072** may buffer (or contain) memory commands (e.g., host memory read or host memory write) directed to the host memory **212** of the host system **620** and other messages directed to the host system **620**. The processing bridge circuit **608N** may include a multiplexer **1011** that variably transfers page miss requests from the host request queue **1074** (e.g., a page response group) and one or more or memory commands or other messages from the outbound queue **1072** into an outbound buffer **1014** of the processing bridge circuit **608N**. The outbound buffer **1014** is thus coupled to the controller memory **1012** and configured to buffer page request groups and the memory commands. In some embodiments, the multiplexer **1011** is a part of outbound queue handler logic **1002** of the processing bridge circuit **608N** that decodes packets from the outbound queue **1072** and the host request queue **1074**.

(129) In embodiments, the queue handling logic **1002** further includes a multiplexer **1016** that determines what content the outbound buffer **1014** contains and directs that content to the correct destination within the processing bridge circuit **608N**. For example, if the outbound buffer **1014** contains a first page request group, the multiplexer **1016** causes the page miss requests of the first page request group to be sent to a translation agent (e.g., TA **617**) of the host system **620**, which causes the translation agent to re-pin a physical page of memory to each respective virtual address of respective page miss requests.

(130) If, however, the queue handling logic **1002** determines the outbound buffer **1014** contains a memory command (such as a read or write directed at the host memory **212**), the multiplexer **1016** directs the memory command into additional outbound queue handling logic **1002**. In various embodiments, the queue handling logic **1002** requests, from the address translation circuit **613** (or ATS), an address translation for a virtual address of the memory command. In response to receiving

the address translation from the address translation circuit **613**, the queue handling logic **1002** sends the memory command with the address translation to the host system **620**, as normal. In response to receiving a miss response from the address translation circuit, however, the queue handling logic **1002** discards the memory command and submits a page miss message to the bridge translation miss queue **1075**. The PRI handler **219** may then process this page miss message along with page miss messages received through the set of translation miss queues **1070**, and submit a corresponding page miss request (or page request group) into the host request queue **1074** for the translation miss. Once an address translation comes back with a page request response from the TA **617**, e.g., via an inbound buffer **1054** (which will be discussed in more depth), the PRI handler **219** can resubmit the command packet (now including an address translation for the virtual address) to the outbound queue **1072**, e.g., an outbound command queue.

(131) In various embodiments, the queue translation logic **1002** buffers read commands into a pending completion tracking queue **1060** that are also sent to the host system **620**. The write commands and the page miss requests may be sent to the host system **620** via an AXI converter **1028**. The AXI converter **1028** can place the write commands and page miss requests out on a set of data channels for writes **1038** to the TA **617** at the host system **620**.

(132) In these embodiments, the read commands may be sent to the host system **620** via an AXI converter **1022**, which sends out read commands to the TA **617** of the host system **620** over a set of address channels for reads **1032**. Further, in these embodiments, the pending completion tracking queue **1060** tracks incoming data received over a set of data channels for reads **1034** and through an AXI converter **1024** that passes the incoming data for the read commands. In embodiments, the pending completion tracking queue **1060** detects data received for a particular read command that is being tracked and, once the data is complete, sends the data to an inbound completion buffer **1064** of the processing bridge circuit **608N**. In at least one embodiment, the inbound completion buffer **1064** receives data of a particular read command from the pending completion tracking queue **1060** and provides the data to an inbound completion queue **1080** stored in the controller memory **1012**.

(133) In various embodiments, page request responses being returned from the TA **617** of the host system **620** are received over a set of data channels for writes **1036** into an AXI converter **1026**. The AXI converter **1026** passes the page request responses (as well as other memory commands and messages, including invalidation requests) received from the host system **620**. The processing bridge circuit **608N** further includes an inbound buffer **1054** to receive the page request responses as well as other memory commands and messages from the host system **620**. In embodiments, inbound queue handling logic **1052** is coupled between the inbound buffer **1054** and the controller memory **1012**.

(134) The inbound queue handling logic **1052** can include a multiplexer **1056** to determine whether the inbound buffer **1054** includes page request responses, memory commands (directed at the controller memory **1012**), messages, or invalidation requests. If invalidation requests, the multiplexer **1056** sends the invalidation requests to the address translation circuit **613**. The multiplexer **1056** sends messages, memory commands, or page request responses to a second multiplexer **1058** that further directs this content. The multiplexer **1058** may send messages and the memory commands to an inbound queue **1082** and the page request responses to a host response queue **1082** of the memory controller **1012**. The host response queue **1084** can buffer the page request responses received from the inbound buffer **1054**. The inbound queue **1082** can buffer incoming memory commands and other messages received from the inbound buffer **1054**.

(135) In at least some embodiments, the PRI handler **219** receives, via the host response queue **1084**, these page request responses corresponding to the page miss requests. In some embodiments, each page request response indicates that the TA **617** has re-pinned a physical page of memory to the virtual address of a respective page miss request. The page request response can further contain an address translation for the virtual address in the original page miss request (or page request

group)), which can be provided to the address translation circuit **613** to be stored in the ATC **616**. (136) In various embodiments, each page request response can be a TLP message sent by the host system **620** to indicate re-pinning failure or success for a group of pages or for some individual pages. It is not an error for the TA **617** of the host system **620** to make a subset of the pages corresponding to the page miss requests resident instead of all the pages. As a result of functionality of the PRI handler **219** in connection with the processing bridge circuit **608N**, the PRI handler **219** may not be allowed to send an “Untranslated Addresses” message even after a successful page request response.

(137) In disclosed embodiments, the controller memory **1012** includes a set of restart queues **1090** that buffer restart messages from the PRI controller **219** corresponding to the page request responses buffered in the host response queue **1084**. In embodiments, each restart queue **1090** is coupled to a respective host interface circuit **608** to provide corresponding restart messages to the respective host interface circuits **608**. More specifically, the host interface circuit **608A** can be coupled to a HIF1 circuit restart queue **1090A**, the host interface circuit **608B** can be coupled to a HIF2 circuit restart queue **1090B**, the host interface circuit **608C** can be coupled to a HIF3 circuit restart queue **1090C**, and a further host interface circuit can be coupled to a HIF4 circuit restart queue **1090D**. Each restart message placed by the PRI handler **219** in one of the set of restart queues **1090** can be provided out of the controller memory **1012** to a corresponding host interface circuit **608**. As discussed with reference to FIG. 9, a restart message causes the host interface circuit **608** to reinsert the command packet into the hardware pipeline **907** where the command packet had originally been removed.

(138) FIG. 11 is a flow chart of an example method **1100** of page request interface support in caching host memory address translation data in a memory sub-system in accordance with some embodiments. The method **1100** can be performed by processing logic that can include hardware (e.g., processing device, circuitry, dedicated logic, programmable logic, microcode, hardware of a device, integrated circuit, etc.), software (e.g., instructions run or executed on a processing device), or a combination thereof. In some embodiments, the method **1100** is performed by the PRI handler **219** of FIG. 2, FIG. 6, and FIG. 9. Although shown in a particular sequence or order, unless otherwise specified, the order of the processes can be modified. Thus, the illustrated embodiments should be understood only as examples, and the illustrated processes can be performed in a different order, and some processes can be performed in parallel. Additionally, one or more processes can be omitted in various embodiments. Thus, not all processes are required in every embodiment. Other process flows are possible.

(139) At operation **1110**, the processing logic stores, in a cache, multiple address translations associated with address translation requests from a host interface circuits. Each address translation request includes a virtual address needing translation.

(140) At operation **1120**, the processing logic tracks translation miss messages received from the host interface circuits, each translation miss message comprising the virtual address of a miss at the cache.

(141) At operation **1130**, the processing logic removes duplicate translation miss messages having an identical virtual address.

(142) At operation **1140**, the processing logic creates multiple page miss requests from non-duplicate translation miss messages that are categorized into page request groups, each page request group corresponding to a host interface circuit of the plurality of host interface circuits.

(143) At operation **1150**, the processing logic queues the page request groups to be sent to a translation agent of a host system. In embodiments, the processing bridge circuit **608N** then sends the queued page request groups to the host system for handling.

(144) FIG. 12 illustrates an example machine of a computer system **1200** within which a set of instructions, for causing the machine to perform any one or more of the methodologies discussed herein, can be executed. In some embodiments, the computer system **1200** can correspond to a host

system (e.g., the host system **120** of FIG. **1**) that includes, is coupled to, or utilizes a memory sub-system (e.g., the memory sub-system **110** of FIG. **1**) or can be used to perform the operations of a controller (e.g., to execute instructions or firmware of the controller **115**). In alternative embodiments, the machine can be connected (e.g., networked) to other machines in a LAN, an intranet, an extranet, and/or the Internet. The machine can operate in the capacity of a server or a client machine in client-server network environment, as a peer machine in a peer-to-peer (or distributed) network environment, or as a server or a client machine in a cloud computing infrastructure or environment.

(145) The machine can be a personal computer (PC), a tablet PC, a set-top box (STB), a Personal Digital Assistant (PDA), a cellular telephone, a web appliance, a server, a network router, a switch or bridge, or any machine capable of executing a set of instructions (sequential or otherwise) that specify actions to be taken by that machine. Further, while a single machine is illustrated, the term “machine” shall also be taken to include any collection of machines that individually or jointly execute a set (or multiple sets) of instructions to perform any one or more of the methodologies discussed herein.

(146) The example computer system **1200** includes a processing device **1202**, a main memory **1204** (e.g., read-only memory (ROM), flash memory, dynamic random access memory (DRAM) such as synchronous DRAM (SDRAM) or Rambus DRAM (RDRAM), etc.), a static memory **1206** (e.g., flash memory, static random access memory (SRAM), etc.), and a data storage system **1218**, which communicate with each other via a bus **1230**.

(147) Processing device **1202** represents one or more general-purpose processing devices such as a microprocessor, a central processing unit, or the like. More particularly, the processing device can be a complex instruction set computing (CISC) microprocessor, reduced instruction set computing (RISC) microprocessor, very long instruction word (VLIW) microprocessor, or a processor implementing other instruction sets, or processors implementing a combination of instruction sets. Processing device **1202** can also be one or more special-purpose processing devices such as an application specific integrated circuit (ASIC), a field programmable gate array (FPGA), a digital signal processor (DSP), network processor, or the like. The processing device **1202** is configured to execute instructions **1226** for performing the operations and steps discussed herein. The computer system **1200** can further include a network interface device **1208** to communicate over the network **1220**.

(148) The data storage system **1218** can include a machine-readable storage medium **1224** (also known as a computer-readable medium) on which is stored one or more sets of instructions **1226** or software embodying any one or more of the methodologies or functions described herein. The instructions **1226** can also reside, completely or at least partially, within the main memory **1204** and/or within the processing device **1202** during execution thereof by the computer system **1200**, the main memory **1204** and the processing device **1202** also constituting machine-readable storage media. The machine-readable storage medium **1224**, data storage system **1218**, and/or main memory **1204** can correspond to the memory sub-system **110** of FIG. **1**.

(149) In one embodiment, the instructions **1226** include instructions to implement functionality corresponding to the controller **115** of FIG. **1**. While the machine-readable storage medium **1224** is shown in an example embodiment to be a single medium, the term “machine-readable storage medium” should be taken to include a single medium or multiple media that store the one or more sets of instructions. The term “machine-readable storage medium” shall also be taken to include any medium that is capable of storing or encoding a set of instructions for execution by the machine and that cause the machine to perform any one or more of the methodologies of the present disclosure. The term “machine-readable storage medium” shall accordingly be taken to include, but not be limited to, solid-state memories, optical media, and magnetic media.

(150) Some portions of the preceding detailed descriptions have been presented in terms of algorithms and symbolic representations of operations on data bits within a computer memory.

These algorithmic descriptions and representations are the ways used by those skilled in the data processing arts to most effectively convey the substance of their work to others skilled in the art. An algorithm is here, and generally, conceived to be a self-consistent sequence of operations leading to a desired result. The operations are those requiring physical manipulations of physical quantities. Usually, though not necessarily, these quantities take the form of electrical or magnetic signals capable of being stored, combined, compared, and otherwise manipulated. It has proven convenient at times, principally for reasons of common usage, to refer to these signals as bits, values, elements, symbols, characters, terms, numbers, or the like.

(151) It should be borne in mind, however, that all of these and similar terms are to be associated with the appropriate physical quantities and are merely convenient labels applied to these quantities. The present disclosure can refer to the action and processes of a computer system, or similar electronic computing device, that manipulates and transforms data represented as physical (electronic) quantities within the computer system's registers and memories into other data similarly represented as physical quantities within the computer system memories or registers or other such information storage systems.

(152) The present disclosure also relates to an apparatus for performing the operations herein. This apparatus can be specially constructed for the intended purposes, or it can include a general purpose computer selectively activated or reconfigured by a computer program stored in the computer. Such a computer program can be stored in a computer readable storage medium, such as, but not limited to, any type of disk including floppy disks, optical disks, CD-ROMs, and magnetic-optical disks, read-only memories (ROMs), random access memories (RAMs), EPROMs, EEPROMs, magnetic or optical cards, or any type of media suitable for storing electronic instructions, each coupled to a computer system bus.

(153) The algorithms and displays presented herein are not inherently related to any particular computer or other apparatus. Various general purpose systems can be used with programs in accordance with the teachings herein, or it can prove convenient to construct a more specialized apparatus to perform the method. The structure for a variety of these systems will appear as set forth in the description below. In addition, the present disclosure is not described with reference to any particular programming language. It will be appreciated that a variety of programming languages can be used to implement the teachings of the disclosure as described herein.

(154) The present disclosure can be provided as a computer program product, or software, that can include a machine-readable medium having stored thereon instructions, which can be used to program a computer system (or other electronic devices) to perform a process according to the present disclosure. A machine-readable medium includes any mechanism for storing information in a form readable by a machine (e.g., a computer). In some embodiments, a machine-readable (e.g., computer-readable) medium includes a machine (e.g., a computer) readable storage medium such as a read only memory ("ROM"), random access memory ("RAM"), magnetic disk storage media, optical storage media, flash memory components, etc.

(155) In the foregoing specification, embodiments of the disclosure have been described with reference to specific example embodiments thereof. It will be evident that various modifications can be made thereto without departing from the broader spirit and scope of embodiments of the disclosure as set forth in the following claims. The specification and drawings are, accordingly, to be regarded in an illustrative sense rather than a restrictive sense.

Claims

1. A processing device comprising: host interface circuitry to interact with a host system; and an address translation circuit of the host interface circuitry to handle address translation requests to the host system from a plurality of host interface circuits, the address translation circuit comprising a cache to store a plurality of address translations associated with the address translation requests;

and a page request interface (PRI) handler one of coupled to or integrated with the host interface circuitry, the PRI handler to perform operations comprising: tracking translation miss messages received from the plurality of host interface circuits, each translation miss message comprising a virtual address of a miss at the cache; removing duplicate translation miss messages having an identical virtual address; categorizing non-duplicate translation miss messages into page request groups based on a host interface circuit, of the plurality of host interface circuits, with which each non-duplicate translation miss message is associated; creating a plurality of page miss requests from the non-duplicate translation miss messages for each page request group; and queuing, in an outbound buffer, the plurality of page miss requests corresponding to one of the page request groups, wherein the plurality of page miss requests are to be sent to a translation agent of the host system.

2. The processing device of claim 1, wherein each page miss request of the plurality of page miss requests is to cause the translation agent to re-pin a physical page of memory to the virtual address of a respective page miss request.

3. The processing device of claim 1, further comprising the plurality of host interface circuits, wherein each host interface circuit of the plurality of host interface circuits affected by a miss at the cache is to: receive a message from the address translation circuit that an address translation request has missed at the cache; remove a command packet from a hardware pipeline, the command packet associated with a function that relies on an address translation corresponding to the address translation request; and send a translation miss message to the PRI handler, the translation miss message comprising the virtual address.

4. The processing device of claim 3, wherein the translation miss message further comprises information that identifies how to reintroduce the command packet into the hardware pipeline.

5. The processing device of claim 3, wherein the operations further comprise: receiving, from the translation agent, a page request response corresponding to the page miss request, the page request response containing an address translation for the virtual address; and causing a restart message to be sent to each respective host interface circuit, the restart message to cause the command packet to be reintroduced to the hardware pipeline.

6. The processing device of claim 3, wherein a first host interface circuit of the plurality of host interface circuits is a processing bridge circuit, the hardware pipeline of the processing bridge circuit comprises an outbound queue containing memory commands, and wherein the operations further comprise resubmitting the command packet comprising an address translation for the virtual address to the outbound queue.

7. The processing device of claim 1, wherein the operations further comprise restricting the plurality of page miss requests that are outstanding for a particular host interface circuit of the plurality of host interface circuits to a threshold number of page miss requests.

8. The processing device of claim 1, further comprising determining, for each non-duplicate translation miss message, an associated host interface circuit by one of: identifying an incoming translation miss queue from which the non-duplicate translation miss message is retrieved; or identifying a hardware identifier within the non-duplicate translation miss message that is indicative of the host interface circuit.

9. A method comprising: storing, in a cache, a plurality of address translations associated with address translation requests from a plurality of host interface circuits and each comprising a virtual address; tracking, by a page request interface (PRI) handler, translation miss messages received from the plurality of host interface circuits, each translation miss message comprising the virtual address of a miss at the cache; removing duplicate translation miss messages having an identical virtual address; categorizing non-duplicate translation miss messages into page request groups based on a host interface circuit, of the plurality of host interface circuits, with which each non-duplicate translation miss message is associated; creating a plurality of page miss requests from the non-duplicate translation miss messages for each page request group; and queuing, by the PRI

handler, in an outbound buffer, the plurality of page miss requests corresponding to one of the page request groups, wherein the plurality of page miss requests are to be sent to a translation agent of a host system.

10. The method of claim 9, wherein each page miss request of the plurality of page miss requests causes the translation agent to re-pin a physical page of memory to the virtual address of a respective page miss request.

11. The method of claim 9, further comprising: receiving, by one or more of the plurality of host interface circuits, a message from an address translation circuit that an address translation request has missed at the cache; removing a command packet from a hardware pipeline, the command packet associated with a function that relies on an address translation corresponding to the address translation request; and sending a translation miss message to the PRI handler, the translation miss message comprising the virtual address.

12. The method of claim 11, wherein the translation miss message further comprises information that identifies how to reintroduce the command packet into the hardware pipeline.

13. The method of claim 11, further comprising: receiving, by the PRI handler, from the translation agent, a page request response corresponding to the page miss request, the page request response containing an address translation for the virtual address; and causing, by the PRI handler, a restart message to be sent to each respective host interface circuit, the restart message to cause the command packet to be reintroduced to the hardware pipeline.

14. The method of claim 11, wherein a first host interface circuit of the plurality of host interface circuits is a processing bridge circuit, the hardware pipeline of the processing bridge circuit comprises an outbound queue containing memory commands, the method further comprise resubmitting, by the PRI handler, the command packet comprising an address translation for the virtual address to the outbound queue.

15. The method of claim 9, further comprising restricting the plurality of page miss requests that are outstanding for a particular host interface circuit of the plurality of host interface circuits to a threshold number of page miss requests.

16. A non-transitory computer-readable storage medium storing instructions, which when executed by a processing device, causes the processing to perform operations comprising: causing to be stored, in a cache, a plurality of address translations associated with address translation requests from a plurality of host interface circuits and each comprising a virtual address; tracking, by a page request interface (PRI) handler, translation miss messages received from the plurality of host interface circuits, each translation miss message comprising the virtual address of a miss at the cache; removing duplicate translation miss messages having an identical virtual address; categorizing non-duplicate translation miss messages into page request groups based on a host interface circuit, of the plurality of host interface circuits, with which each non-duplicate translation miss message is associated; creating a plurality of page miss requests from the non-duplicate translation miss messages for each page request group; and queuing, by the PRI handler, in an outbound buffer, the plurality of page miss requests corresponding to one of the page request groups, wherein the plurality of page miss requests are to be sent to a translation agent of a host system.

17. The non-transitory computer-readable storage medium of claim 16, wherein each page miss request of the plurality of page miss requests causes the translation agent to re-pin a physical page of memory to the virtual address of a respective page miss request.

18. The non-transitory computer-readable storage medium of claim 16, wherein the operations further comprise: receiving, by one or more of the plurality of host interface circuits, a message from an address translation circuit that an address translation request has missed at the cache; removing a command packet from a hardware pipeline, the command packet associated with a function that relies on an address translation corresponding to the address translation request; and sending a translation miss message to the PRI handler, the translation miss message comprising the

virtual address.

19. The non-transitory computer-readable storage medium of claim 18, wherein the translation miss message further comprises information that identifies how to reintroduce the command packet into the hardware pipeline, and wherein the operations further comprise: receiving, by the PRI handler, from the translation agent, a page request response corresponding to the page miss request, the page request response containing an address translation for the virtual address; and causing, by the PRI handler, a restart message to be sent to each respective host interface circuit, the restart message to cause the command packet to be reintroduced to the hardware pipeline.

20. The non-transitory computer-readable storage medium of claim 18, wherein a first host interface circuit of the plurality of host interface circuits is a processing bridge circuit, the hardware pipeline of the processing bridge circuit comprises an outbound queue containing memory commands, and the operations further comprise resubmitting, by the PRI handler, the command packet comprising an address translation for the virtual address to the outbound queue.

21. The non-transitory computer-readable storage medium of claim 16, wherein the operations further comprise restricting the plurality of page miss requests that are outstanding for a particular host interface circuit of the plurality of host interface circuits to a threshold number of page miss requests.
