

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417386
Kind Code	B2
Date of Patent	September 16, 2025
Inventor(s)	Galvin; Brian

System and methods for upsampling of decompressed genomic data after lossy compression using a neural network

Abstract

A system and methods for upsampling of decompressed genomic data after lossy compression using a neural network integrates AI-based techniques to enhance compression quality. It incorporates a novel deep-learning neural network that upsamples decompressed data to restore information lost during lossy compression, taking advantage of cross-correlations between genomic data sets.

Inventors:	Galvin; Brian (Silverdale, WA)
Applicant:	AtomBeam Technologies Inc. (Moraga, CA)
Family ID:	1000008822156
Assignee:	ATOMBEAM TECHNOLOGIES INC. (Moraga, CA)
Appl. No.:	19/007525
Filed:	January 01, 2025

Prior Publication Data

Document Identifier	Publication Date
US 20250190790 A1	Jun. 12, 2025

Related U.S. Application Data

continuation-in-part parent-doc US 18737906 20240607 PENDING child-doc US 19007525
continuation-in-part parent-doc US 18420771 20240124 US 12095484 20240917 child-doc US 19007525
continuation-in-part parent-doc US 18410980 20240111 US 12068761 20240820 child-doc US 18737906

Publication Classification

Int. Cl.: H03M7/00 (20060101); G06N3/08 (20230101)

U.S. Cl.:

CPC G06N3/08 (20130101);

Field of Classification Search

CPC: H03M (7/3062); H04N (19/80)

USPC: 341/50; 341/51

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2018/0300606	12/2017	Corkery	N/A	G06F 12/0862

Primary Examiner: Jeanglaude; Jean B

Attorney, Agent or Firm: Galvin Patent Law LLC

Background/Summary

CROSS-REFERENCE TO RELATED APPLICATIONS

(1) Priority is claimed in the application data sheet to the following patents or patent applications, each of which is expressly incorporated herein by reference in its entirety: Ser. No. 18/737,906 Ser. No. 18/420,771 Ser. No. 18/410,980 Ser. No. 18/537,728

BACKGROUND OF THE INVENTION

Field of the Art

(2) The present invention relates to the fields of bioinformatics and computational biology, specifically to the processing, analysis, and integration of genomic and proteomic data using advanced machine learning techniques, with a focus on lossy compression and subsequent upsampling of omics data.

Discussion of the State of the Art

(3) For many applications, such as video compression for streaming video, lossy compression techniques such as HEVC (high-efficiency video coding) to optimize the use of available bandwidth and for other purposes. By definition, lossy compression involves the loss of some of the data being transmitted in the process of compression; in the video compression example, this results in lower-resolution video and provides the reason for pixelated video in low-bandwidth situations. Clearly it would be desirable to recover as much of the lost data as possible, but of course this is impossible in a single compressed channel for the method of compression results in a true loss of information.

(4) Genomic data is highly diverse and complex. The information is encoded in DNA sequences, which can vary significantly among individuals and species. Creating a compression algorithm that effectively handles this diversity can be challenging. Balancing between lossless and lossy compression is a challenge. While lossless compression retains all the original information, it may not achieve high compression ratios. On the other hand, lossy compression may achieve higher ratios but at the cost of losing some information, which may be critical in genomics. As the volume of genomic data continues to increase with advances in sequencing technologies, the scalability of compression algorithms becomes crucial. Some algorithms may not scale well to handle the growing amount of data. Compressed genomic data should still be interpretable and accessible by researchers and clinicians. Ensuring that the compressed data retains its biological significance is important for downstream analyses.

(5) In recent years, high-throughput technologies in genomics and proteomics have led to an explosion of biological data. Next-generation sequencing has dramatically reduced the cost and time required for genomic analysis, while mass spectrometry-based techniques have enabled large-scale proteomic studies. However, the sheer volume and complexity of this data present significant challenges for storage, processing, and interpretation.

(6) Genomic data compression has become crucial for efficient storage and transmission. Existing compression methods often struggle to balance high compression ratios with the need to preserve biologically relevant information. Moreover, lossy compression techniques, while achieving higher compression rates, risk losing critical genetic variations that could be important for downstream analyses.

(7) In proteomics, the challenges are even more pronounced. Proteomic data is inherently more complex than genomic data due to the dynamic nature of proteins, post-translational modifications, and the importance of three-dimensional structure in protein function. Current proteomic data analysis pipelines face several hurdles. Mass spectrometry experiments generate vast amounts of complex spectral data that require sophisticated algorithms for interpretation. Accurately identifying and quantifying proteins from mass spectrometry data remains challenging, especially for low-abundance proteins or complex mixtures. While protein sequence data is more readily available, predicting three-dimensional protein structures—crucial for understanding function—remains computationally intensive and often inaccurate for novel proteins. Identifying potential binding sites on proteins is essential for understanding protein-protein interactions and drug design, but current methods often lack accuracy or are computationally prohibitive for large-scale analyses. Integrating diverse types of proteomic data (e.g., abundance, post-translational modifications, structural information) with genomic data and other omics data types remains a significant challenge.

(8) Furthermore, the post-processing of both genomic and proteomic data is critical for ensuring the quality and reliability of analysis results. Current post-processing methods often lack standardization and fail to adequately address issues such as false discovery rates in protein identification, accurate quantification in the presence of missing data, and the integration of multiple data types for comprehensive biological insights.

(9) Existing machine learning approaches, particularly deep learning models, have shown promise in addressing some of these challenges. However, most current models are designed to handle either genomic or proteomic data, not both. Additionally, these models often require extensive computational resources and struggle with the interpretability of their results, a crucial factor in biological research.

(10) There is a pressing need for a unified system that can efficiently compress, process, analyze, and integrate both genomic and proteomic data. Such a system should be capable of handling the high dimensionality and complexity of omics data, provide accurate predictions of protein structures and binding sites, and offer interpretable results that can drive biological discovery. Moreover, it should incorporate robust post-processing techniques to ensure the reliability and

biological relevance of the outputs.

(11) What is needed is a novel approach that combines advanced machine learning techniques, including neural networks and graph-based representations, with domain-specific biological knowledge to create a comprehensive platform for genomic and proteomic data analysis. This approach should address the challenges of data compression, upsampling, structural prediction, and data integration, while also providing sophisticated post-processing capabilities to refine and validate the results.

(12) What is needed is a system and methods for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks, with integrated structural prediction and post-processing capabilities.

SUMMARY OF THE INVENTION

(13) Accordingly, the inventor has conceived and reduced to practice, a system and methods for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks that integrate AI-based techniques to enhance compression quality and facilitate integrated analysis. It incorporates novel neural network architectures including ProteomicUNet for upsampling, ProteinFormer for 3D structure prediction, and BindingSiteFinder for binding site prediction. Additionally, the system employs a deep learning system with a latent transformer core for large codeword models, enabling efficient processing of extensive genomic and proteomic datasets while maintaining high fidelity in data representation. These components work in concert to address both local and global features, mitigating compression artifacts and improving decompressed data quality while enabling advanced genomic and proteomic analysis. The system's outputs enable effective data reconstruction and integrated analysis, achieving advanced compression while preserving crucial information for accurate genomic and proteomic investigation.

(14) According to a preferred embodiment, a system for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks is disclosed, comprising: a computing system comprising at least a memory and a processor; two or more datasets that are substantially correlated and which have been compressed with lossy compression, the two or more datasets comprising genomic data and proteomic data; a deep learning neural network configured to recover lost information associated with a compressed bit stream; a decoder configured to receive and decompress a compressed bit stream comprising cross-correlated genomic data and proteomic data; a proteomic data upsampling neural network configured to upsample compressed proteomic data; a protein structure prediction model configured to predict 3D protein structures from upsampled proteomic data; a binding site prediction model configured to predict binding sites using the predicted 3D protein structures and upsampled proteomic data; and a deep learning system with a latent transformer core for large codeword models, configured to process input vectors through a variational autoencoder and a transformer to generate output vectors.

(15) According to another preferred embodiment, a method for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks is disclosed, comprising the steps of: receiving a compressed bit stream comprising cross-correlated genomic data and proteomic data; decompressing the compressed bit stream; using the decompressed bit stream as input into a deep learning neural network to recover information lost during lossy compression; upsampling compressed proteomic data; predicting 3D protein structures from the upsampled proteomic data; predicting binding sites using the predicted 3D protein structures and upsampled proteomic data; and processing input vectors through a deep learning system with a latent transformer core for large codeword models to generate output vectors.

(16) According to an aspect of an embodiment, the genomic data comprises parallel genome datasets.

(17) According to an aspect of an embodiment, the two or more datasets comprise genomic data from a subset of the human genome.

- (18) According to an aspect of an embodiment, the deep learning neural network is a neural network that can recover signals from a compressed bitstream.
- (19) According to an aspect of an embodiment, the compressed bit stream comprises a plurality of channels, wherein each of the plurality of channels is associated with a genomic dataset or a proteomic dataset.
- (20) According to an aspect of an embodiment, the system further comprises a multi-modal protein data graph representation generator configured to create a multi-modal graph representation integrating proteomic data, predicted structures, and binding sites.
- (21) According to an aspect of an embodiment, the input vectors processed by the deep learning system with a latent transformer core may contain appended zeros, truncated data points, or appended metadata.
-

Description

BRIEF DESCRIPTION OF THE DRAWING FIGURES

- (1) FIG. 1 is a block diagram illustrating an exemplary system architecture for upsampling of decompressed data after lossy compression using a neural network, according to an embodiment.
- (2) FIGS. 2A and 2B illustrate an exemplary architecture for an AI deblocking network configured to provide deblocking on dual-channel data stream comprising SAR I/Q data, according to an embodiment.
- (3) FIG. 3 is a block diagram illustrating an exemplary architecture for a component of the system for SAR image compression, the channel-wise transformer.
- (4) FIG. 4 is a block diagram illustrating an exemplary system architecture for providing lossless data compaction, according to an embodiment.
- (5) FIG. 5 is a diagram showing an embodiment of one aspect of the lossless data compaction system, specifically data deconstruction engine.
- (6) FIG. 6 is a diagram showing an embodiment of another aspect of lossless data compaction system, specifically data reconstruction engine.
- (7) FIG. 7 is a diagram showing an embodiment of another aspect of lossless data compaction the system, specifically library manager.
- (8) FIG. 8 is a flow diagram illustrating an exemplary method for complex-valued SAR image compression, according to an embodiment.
- (9) FIG. 9 is a flow diagram illustrating an exemplary method for decompression of a complex-valued SAR image, according to an embodiment.
- (10) FIG. 10 is a flow diagram illustrating an exemplary method for deblocking using a trained deep learning algorithm, according to an embodiment.
- (11) FIGS. 11A and 11B illustrate an exemplary architecture for an AI deblocking network configured to provide deblocking for a general N-channel data stream, according to an embodiment.
- (12) FIG. 12 is a block diagram illustrating an exemplary system architecture for N-channel data compression with predictive recovery, according to an embodiment.
- (13) FIG. 13 is a flow diagram illustrating an exemplary method for processing a compressed n-channel bit stream using an AI deblocking network, according to an embodiment.
- (14) FIG. 14 is a block diagram illustrating a system for training a neural network to perform upsampling of decompressed data after lossy compression, according to an embodiment.
- (15) FIG. 15 is a flow diagram illustrating an exemplary method for training a neural network to perform upsampling of decompressed data after lossy compression, according to an embodiment.
- (16) FIG. 16 is a block diagram illustrating an exemplary architecture for a neural upsampler configured to process N-channel time-series data, according to an embodiment.

- (17) FIG. 17 is a block diagram illustrating an exemplary system architecture for upsampling of decompressed sensor data after lossy compression using a neural network, according to an embodiment.
- (18) FIG. 18 is a flow diagram illustrating an exemplary method for performing neural upsampling of two or more time-series data streams, according to an embodiment.
- (19) FIG. 19 is a block diagram illustrating an exemplary system architecture for neural upsampling of two or more genomic datasets, according to an embodiment.
- (20) FIG. 20 is a flow diagram illustrating an exemplary method for performing neural upsampling of two or more genomic datasets, according to an embodiment.
- (21) FIG. 21 illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part.
- (22) FIG. 22 is a block diagram illustrating an exemplary architecture for ProteomicUNet, a neural network configured to upsample compressed proteomic data.
- (23) FIG. 23 is a method diagram illustrating an exemplary method for training ProteomicUNet.
- (24) FIG. 24 is a method diagram illustrating an exemplary method for utilizing the trained
- (25) ProteomicUNet.
- (26) FIG. 25 is a block diagram illustrating an exemplary architecture for ProteinFormer for protein structure prediction.
- (27) FIG. 26 is a method diagram illustrating an exemplary method for utilizing ProteinFormer for protein structure prediction
- (28) FIG. 27 is a block diagram illustrating an exemplary architecture for BindingSiteFinder to predict protein binding sites.
- (29) FIG. 28 is a method diagram illustrating an exemplary method for utilizing BindingSiteFinder to predict protein binding sites.
- (30) FIG. 29 is a block diagram illustrating an exemplary integrated system architecture combining ProteinFormer and BindingSiteFinder for comprehensive protein structure and binding site prediction.
- (31) FIG. 30 is a method diagram illustrating an exemplary method for the integrated use of ProteinFormer and BindingSiteFinder for comprehensive protein structure and binding site prediction.
- (32) FIG. 31 is a block diagram illustrating the structure of ProteoGraph, a multi-modal graph representation designed to integrate various types of proteomic data.
- (33) FIG. 32 is a method diagram illustrating an exemplary method for constructing and utilizing ProteoGraph.
- (34) FIG. 33 is a block diagram illustrating the architecture of the deep learning system with latent transformer core.
- (35) FIG. 34 is a method diagram illustrating the use of latent space transformation for genomic and proteomic data.
- (36) FIG. 35 is a block diagram illustrating data flow through components of the system, in an embodiment.

DETAILED DESCRIPTION OF THE INVENTION

(37) The inventor has conceived, and reduced to practice, a system and methods for upsampling of decompressed genomic data after lossy compression using a neural network that integrates AI-based techniques to enhance compression quality. It incorporates a novel AI deblocking network composed of recurrent layers for feature extraction and a channel-wise transformer with attention to capture complex inter-channel dependencies. The recurrent layers extract multi-dimensional features from the two or more correlated datasets, while the channel-wise transformer learns global inter-channel relationships. This hybrid approach addresses both local and global features, mitigating compression artifacts and improving decompressed data quality. The model's outputs enable effective data reconstruction, achieving advanced compression while preserving crucial

information for accurate analysis.

(38) SAR images provide an excellent exemplary use case for a system and methods for upsampling of decompressed data after lossy compression. Synthetic Aperture Radar technology is used to capture detailed images of the Earth's surface by emitting microwave signals and measuring their reflections. Unlike traditional grayscale images that use a single intensity value per pixel, SAR images are more complex. Each pixel in a SAR image contains not just one value but a complex number ($I+Qi$). A complex number consists of two components: magnitude (or amplitude) and phase. In the context of SAR, the complex value at each pixel represents the strength of the radar signal's reflection (magnitude) and the phase shift (phase) of the signal after interacting with the terrain. This information is crucial for understanding the properties of the surface and the objects present. In a complex-value SAR image, the magnitude of the complex number indicates the intensity of the radar reflection, essentially representing how strong the radar signal bounced back from the surface. Higher magnitudes usually correspond to stronger reflections, which may indicate dense or reflective materials on the ground.

(39) The complex nature of SAR images stems from the interference and coherence properties of radar waves. When radar waves bounce off various features on the Earth's surface, they can interfere with each other. This interference pattern depends on the radar's wavelength, the angle of incidence, and the distances the waves travel. As a result, the radar waves can combine constructively (amplifying the signal) or destructively (canceling out the signal). This interference phenomenon contributes to the complex nature of SAR images. The phase of the complex value encodes information about the distance the radar signal traveled and any changes it underwent during the round-trip journey. For instance, if the radar signal encounters a surface that's slightly elevated or depressed, the phase of the returning signal will be shifted accordingly. Phase information is crucial for generating accurate topographic maps and understanding the geometry of the terrain.

(40) Coherence refers to the consistency of the phase relationship between different pixels in a SAR image. Regions with high coherence have similar phase patterns and are likely to represent stable surfaces or structures, while regions with low coherence might indicate changes or disturbances in the terrain.

(41) Complex-value SAR image compression is important for several reasons such as data volume reduction, bandwidth and transmission efficiency, real-time applications, and archiving and retrieval. SAR images can be quite large due to their high resolution and complex nature. Compression helps reduce the storage and transmission requirements, making it more feasible to handle and process the data. When SAR images need to be transmitted over limited bandwidth channels, compression can help optimize data transmission and minimize communication costs. Some SAR applications, such as disaster response and surveillance, require real-time processing. Compressed data can be processed faster, enabling quicker decision-making. Additionally, compressed SAR images take up less storage space, making long-term archiving and retrieval more manageable.

(42) According to various embodiments, a system is proposed which provides a novel pipeline for compressing and subsequently recovering complex-valued SAR image data using a prediction recovery framework that utilizes a conventional image compression algorithm to encode the original image to a bitstream. In an embodiment, a lossless compaction method may be applied to the encoded bitstream, further reducing the size of the SAR image data for both storage and transmission. Subsequently, the system decodes a prediction of the I/Q channels and then recovers the phase and amplitude via a deep-learning based network to effectively remove compression artifacts and recover information of the SAR image as part of the loss function in the training. The deep-learning based network may be referred to herein as an artificial intelligence (AI) deblocking network.

(43) Deblocking refers to a technique used to reduce or eliminate blocky artifacts that can occur in

compressed images or videos. These artifacts are a result of lossy compression algorithms, such as JPEG for images or various video codecs like H.264, H.265 (HEVC), and others, which divide the image or video into blocks and encode them with varying levels of quality. Blocky artifacts, also known as “blocking artifacts,” become visible when the compression ratio is high, or the bitrate is low. These artifacts manifest as noticeable edges or discontinuities between adjacent blocks in the image or video. The result is a visual degradation characterized by visible square or rectangular regions, which can significantly reduce the overall quality and aesthetics of the content. Deblocking techniques are applied during the decoding process to mitigate or remove these artifacts. These techniques typically involve post-processing steps that smooth out the transitions between adjacent blocks, thus improving the overall visual appearance of the image or video. Deblocking filters are commonly used in video codecs to reduce the impact of blocking artifacts on the decoded video frames.

(44) According to various embodiments, the disclosed system and methods may utilize a SAR recovery network configured to perform data deblocking during the data decoding process. Amplitude and phase images exhibit a non-linear relationship, while I and Q images demonstrate a linear relationship. The SAR recovery network is designed to leverage this linear relationship by utilizing the I/Q images to enhance the decoded SAR image. In an embodiment, the SAR recovery network is a deep learned neural network. According to an aspect of an embodiment, the SAR recovery network utilizes residual learning techniques. According to an aspect of an embodiment, the SAR recovery network comprises a channel-wise transformer with attention. According to an aspect of an embodiment, the SAR recovery network comprises Multi-Scale Attention Blocks (MSAB).

(45) A channel-wise transformer with attention is a neural network architecture that combines elements of both the transformer architecture and channel-wise attention mechanisms. It's designed to process multi-channel data, such as SAR images, where each channel corresponds to a specific feature map or modality. The transformer architecture is a powerful neural network architecture initially designed for natural language processing (NLP) tasks. It consists of self-attention mechanisms that allow each element in a sequence to capture relationships with other elements, regardless of their position. The transformer has two main components: the self-attention mechanism (multi-head self-attention) and feedforward neural networks (position-wise feedforward layers). Channel-wise attention, also known as “Squeeze-and-Excitation” (SE) attention, is a mechanism commonly used in convolutional neural networks (CNNs) to model the interdependencies between channels (feature maps) within a single layer. It assigns different weights to different channels to emphasize important channels and suppress less informative ones. At each layer of the network, a channel-wise attention mechanism is applied to the input data. This mechanism captures the relationships between different channels within the same layer and assigns importance scores to each channel based on its contribution to the overall representation. After the channel-wise attention, a transformer-style self-attention mechanism is applied to the output of the channel-wise attention. This allows each channel to capture dependencies with other channels in a more global context, similar to how the transformer captures relationships between elements in a sequence. Following the transformer self-attention, feedforward neural network layers (position-wise feedforward layers) can be applied to further process the transformed data.

(46) According to various embodiments, the system incorporates a novel pipeline for processing and analyzing proteomic data using advanced neural network architectures. This pipeline includes components for upsampling compressed proteomic data, predicting 3D protein structures, and identifying binding sites. The system utilizes a prediction recovery framework that leverages deep learning techniques to effectively reconstruct high-quality proteomic data from compressed formats and perform advanced analyses.

(47) A key component of this pipeline is the ProteomicUNet, a neural network designed specifically for upsampling compressed proteomic data. ProteomicUNet combines convolutional

layers, LSTM layers, and attention mechanisms to recover information lost during the compression process. The network architecture is tailored to handle the unique characteristics of proteomic data, such as peptide sequences and mass spectrometry patterns. ProteomicUNet effectively mitigates compression artifacts and enhances the quality of the reconstructed proteomic data, enabling more accurate downstream analyses.

(48) Following the upsampling process, the system employs ProteinFormer, a transformer-based model for predicting 3D protein structures from the upsampled proteomic data. ProteinFormer leverages the power of self-attention mechanisms to capture long-range dependencies in protein sequences, which are crucial for accurate structure prediction. The model is trained using a novel self-supervised pre-training task, where it predicts masked amino acids in protein sequences to learn structural patterns. This approach allows ProteinFormer to generate high-quality 3D protein structure predictions, which are essential for understanding protein function and interactions.

(49) The system also incorporates BindingSiteFinder, a graph neural network-based model for predicting protein binding sites. BindingSiteFinder integrates information from the upsampled proteomic data, predicted 3D structures, and evolutionary conservation data. The model employs a novel message-passing scheme that incorporates features from multiple data sources, allowing it to identify potential binding sites with high accuracy. BindingSiteFinder's predictions are crucial for understanding protein-protein interactions and identifying potential drug targets.

(50) To facilitate comprehensive analysis and interpretation of the processed data, the system generates a ProteoGraph representation. ProteoGraph is a multi-modal graph structure that integrates various types of proteomic data and captures complex relationships between them. It represents proteins, peptides, post-translational modifications, predicted structures, and binding sites as nodes, with edges representing the relationships between these entities. This unified representation allows researchers to visualize and analyze complex proteomic data in a coherent framework, enabling deeper insights into biological processes.

(51) The system incorporates a deep learning architecture with a latent transformer core, specifically designed for processing large codeword models in genomic and proteomic data analysis. This component leverages the power of transformer models in a latent space, allowing for efficient processing of extensive datasets while maintaining high fidelity in data representation. The architecture comprises a variational autoencoder (VAE) and a transformer model. Input vectors, which may include genomic or proteomic data, are first processed through the encoder of the VAE to generate latent space vectors. These latent representations are then fed into a transformer model, which, notably, does not include the typical embedding and positional encoding layers. This design choice allows the transformer to focus solely on learning relationships between the latent space vectors. The transformer's output is then used to generate output latent space vectors, which are finally decoded back into the original data space using the VAE's decoder. This approach enables the system to capture complex, long-range dependencies in the data, which is crucial for accurate compression and subsequent analysis of genomic and proteomic information. By operating in the latent space, this architecture can handle much larger and more complex codeword models than traditional approaches, significantly enhancing the system's ability to process and analyze extensive genomic and proteomic datasets with minimal loss of information.

(52) The system and methods described herein in various embodiments may be directed to the processing of audio data such as, for example, speech channels associated with one or more individuals.

(53) One or more different aspects may be described in the present application. Further, for one or more of the aspects described herein, numerous alternative arrangements may be described; it should be appreciated that these are presented for illustrative purposes only and are not limiting of the aspects contained herein or the claims presented herein in any way. One or more of the arrangements may be widely applicable to numerous aspects, as may be readily apparent from the disclosure. In general, arrangements are described in sufficient detail to enable those skilled in the

art to practice one or more of the aspects, and it should be appreciated that other arrangements may be utilized and that structural, logical, software, electrical and other changes may be made without departing from the scope of the particular aspects. Particular features of one or more of the aspects described herein may be described with reference to one or more particular aspects or figures that form a part of the present disclosure, and in which are shown, by way of illustration, specific arrangements of one or more of the aspects. It should be appreciated, however, that such features are not limited to usage in the one or more particular aspects or figures with reference to which they are described. The present disclosure is neither a literal description of all arrangements of one or more of the aspects nor a listing of features of one or more of the aspects that must be present in all arrangements.

(54) Headings of sections provided in this patent application and the title of this patent application are for convenience only, and are not to be taken as limiting the disclosure in any way.

(55) Devices that are in communication with each other need not be in continuous communication with each other, unless expressly specified otherwise. In addition, devices that are in communication with each other may communicate directly or indirectly through one or more communication means or intermediaries, logical or physical.

(56) A description of an aspect with several components in communication with each other does not imply that all such components are required. To the contrary, a variety of optional components may be described to illustrate a wide variety of possible aspects and in order to more fully illustrate one or more aspects. Similarly, although process steps, method steps, algorithms or the like may be described in a sequential order, such processes, methods and algorithms may generally be configured to work in alternate orders, unless specifically stated to the contrary. In other words, any sequence or order of steps that may be described in this patent application does not, in and of itself, indicate a requirement that the steps be performed in that order. The steps of described processes may be performed in any order practical. Further, some steps may be performed simultaneously despite being described or implied as occurring non-simultaneously (e.g., because one step is described after the other step). Moreover, the illustration of a process by its depiction in a drawing does not imply that the illustrated process is exclusive of other variations and modifications thereto, does not imply that the illustrated process or any of its steps are necessary to one or more of the aspects, and does not imply that the illustrated process is preferred. Also, steps are generally described once per aspect, but this does not mean they must occur once, or that they may only occur once each time a process, method, or algorithm is carried out or executed. Some steps may be omitted in some aspects or some occurrences, or some steps may be executed more than once in a given aspect or occurrence.

(57) When a single device or article is described herein, it will be readily apparent that more than one device or article may be used in place of a single device or article. Similarly, where more than one device or article is described herein, it will be readily apparent that a single device or article may be used in place of the more than one device or article.

(58) The functionality or the features of a device may be alternatively embodied by one or more other devices that are not explicitly described as having such functionality or features. Thus, other aspects need not include the device itself.

(59) Techniques and mechanisms described or referenced herein will sometimes be described in singular form for clarity. However, it should be appreciated that particular aspects may include multiple iterations of a technique or multiple instantiations of a mechanism unless noted otherwise. Process descriptions or blocks in figures should be understood as representing modules, segments, or portions of code which include one or more executable instructions for implementing specific logical functions or steps in the process. Alternate implementations are included within the scope of various aspects in which, for example, functions may be executed out of order from that shown or discussed, including substantially concurrently or in reverse order, depending on the functionality involved, as would be understood by those having ordinary skill in the art.

Definitions

(60) The term “bit” refers to the smallest unit of information that can be stored or transmitted. It is in the form of a binary digit (either 0 or 1). In terms of hardware, the bit is represented as an electrical signal that is either off (representing 0) or on (representing 1).

(61) The term “codebook” refers to a database containing sourceblocks each with a pattern of bits and reference code unique within that library. The terms “library” and “encoding/decoding library” are synonymous with the term codebook.

(62) The terms “compression” and “deflation” as used herein mean the representation of data in a more compact form than the original dataset. Compression and/or deflation may be either “lossless”, in which the data can be reconstructed in its original form without any loss of the original data, or “lossy” in which the data can be reconstructed in its original form, but with some loss of the original data.

(63) The terms “compression factor” and “deflation factor” as used herein mean the net reduction in size of the compressed data relative to the original data (e.g., if the new data is 70% of the size of the original, then the deflation/compression factor is 30% or 0.3.)

(64) The terms “compression ratio” and “deflation ratio”, and as used herein all mean the size of the original data relative to the size of the compressed data (e.g., if the new data is 70% of the size of the original, then the deflation/compression ratio is 70% or 0.7.)

(65) The term “data set” refers to a grouping of data for a particular purpose. One example of a data set might be a word processing file containing text and formatting information. Another example of a data set might comprise data gathered/generated as the result of one or more radars in operation.

(66) The term “sourcepacket” as used herein means a packet of data received for encoding or decoding. A sourcepacket may be a portion of a data set.

(67) The term “sourceblock” as used herein means a defined number of bits or bytes used as the block size for encoding or decoding. A sourcepacket may be divisible into a number of sourceblocks. As one non-limiting example, a 1 megabyte sourcepacket of data may be encoded using 512 byte sourceblocks. The number of bits in a sourceblock may be dynamically optimized by the system during operation. In one aspect, a sourceblock may be of the same length as the block size used by a particular file system, typically 512 bytes or 4,096 bytes.

(68) The term “codeword” refers to the reference code form in which data is stored or transmitted in an aspect of the system. A codeword consists of a reference code to a sourceblock in the library plus an indication of that sourceblock's location in a particular data set.

(69) The term “deblocking” as used herein refers to a technique used to reduce or eliminate blocky artifacts that can occur in compressed images or videos. These artifacts are a result of lossy compression algorithms, such as JPEG for images or various video codecs like H.264, H.265 (HEVC), and others, which divide the image or video into blocks and encode them with varying levels of quality. Blocky artifacts, also known as “blocking artifacts,” become visible when the compression ratio is high, or the bitrate is low. These artifacts manifest as noticeable edges or discontinuities between adjacent blocks in the image or video. The result is a visual degradation characterized by visible square or rectangular regions, which can significantly reduce the overall quality and aesthetics of the content. Deblocking techniques are applied during the decoding process to mitigate or remove these artifacts. These techniques typically involve post-processing steps that smooth out the transitions between adjacent blocks, thus improving the overall visual appearance of the image or video. Deblocking filters are commonly used in video codecs to reduce the impact of blocking artifacts on the decoded video frames. A primary goal of deblocking is to enhance the perceptual quality of the compressed content, making it more visually appealing to viewers. It's important to note that deblocking is just one of many post-processing steps applied during the decoding and playback of compressed images and videos to improve their quality.

(70) Conceptual Architecture

(71) FIG. 1 is a block diagram illustrating an exemplary system architecture **100** for upsampling of

decompressed data after lossy compression using a neural network, according to an embodiment. According to the embodiment, the system **100** comprises an encoder module **110** configured to receive two or more datasets **101a-n** which are substantially correlated and perform lossy compression on the received dataset, and a decoder module **120** configured to receive a compressed bit stream and use a trained neural network to output a reconstructed dataset which can restore most of the “lost” data due to the lossy compression. Datasets **101a-n** may comprise streaming data or data received in a batch format. Datasets **101a-n** may comprise one or more datasets, data streams, data files, or various other types of data structures which may be compressed. Furthermore, dataset **101a-n** may comprise n-channel data comprising a plurality of data channels sent via a single data stream.

(72) Encoder **110** may utilize a lossy compression module **111** to perform lossy compression on a received dataset **101a-n**. The type of lossy compression implemented by lossy compression module **111** may be dependent upon the data type being processed. For example, for SAR imagery data, High Efficiency Video Coding (HEVC) may be used to compress the dataset. In another example, if the data being processed is time-series data, then delta encoding may be used to compress the dataset. The encoder **110** may then send the compressed data as a compressed data stream to a decoder **120** which can receive the compressed data stream and decompress the data using a decompression module **121**.

(73) The decompression module **121** may be configured to perform data decompression a compressed data stream using an appropriate data decompression algorithm. The decompressed data may then be used as input to a neural upsampler **122** which utilizes a trained neural network to restore the decompressed data to nearly its original state **105** by taking advantage of the information embedded in the correlation between the two or more datasets **101a-n**.

(74) FIGS. 2A and 2B illustrate an exemplary architecture for an AI deblocking network configured to provide deblocking for dual-channel data stream comprising SAR I/Q data, according to an embodiment. In the context of this disclosure, dual-channel data refers to fact that SAR image signal can be represented as two (dual) components (i.e., I and Q) which are correlated to each other in some manner. In the case of I and Q, their correlation is that they can be transformed into phase and amplitude information and vice versa. AI deblocking network utilizes a deep learned neural network architecture for joint frequency and pixel domain learning. According to the embodiment, a network may be developed for joint learning across one or more domains. As shown, the top branch **210** is associated with the pixel domain learning and the bottom branch **220** is associated with the frequency domain learning. According to the embodiment, the AI deblocking network receives as input complex-valued SAR image I and Q channels **201** which, having been encoded via encoder **110**, has subsequently been decompressed via decoder **120** before being passed to AI deblocking network for image enhancement via artifact removal. Inspired by the residual learning network and the MSAB attention mechanism, AI deblocking network employs resblocks that take two inputs. In some implementations, to reduce complexity the spatial resolution may be downsampled to one-half and one-fourth. During the final reconstruction the data may be upsampled to its original resolution. In one implementation, in addition to downsampling, the network employs deformable convolution to extract initial features, which are then passed to the resblocks. In an embodiment, the network comprises one or more resblocks and one or more convolutional filters. In an embodiment, the network comprises 8 resblocks and 64 convolutional filters.

(75) Deformable convolution is a type of convolutional operation that introduces spatial deformations to the standard convolutional grid, allowing the convolutional kernel to adaptively sample input features based on the learned offsets. It's a technique designed to enhance the modeling of spatial relationships and adapt to object deformations in computer vision tasks. In traditional convolutional operations, the kernel's positions are fixed and aligned on a regular grid across the input feature map. This fixed grid can limit the ability of the convolutional layer to

capture complex transformations, non-rigid deformations, and variations in object appearance. Deformable convolution aims to address this limitation by introducing the concept of spatial deformations. Deformable convolution has been particularly effective in tasks like object detection and semantic segmentation, where capturing object deformations and accurately localizing object boundaries are important. By allowing the convolutional kernels to adaptively sample input features from different positions based on learned offsets, deformable convolution can improve the model's ability to handle complex and diverse visual patterns.

(76) According to an embodiment, the network may be trained as a two stage process, each utilizing specific loss functions. During the first stage, a mean squared error (MSE) function is used in the I/Q domain as a primary loss function for the AI deblocking network. The loss function of the SAR I/Q channel LSAR is defined as:

$$L_{\text{sub.SAR}} = E[\|I - I_{\text{sub.amp}}\|_{\text{sub.2}}^2]$$

(77) Moving to the second stage, the network reconstructs the amplitude component and computes the amplitude loss using MSE as follows:

$$L_{\text{sub.amp}} = E[\|I_{\text{sub.amp}} - I_{\text{sub.dec,amp}}\|_{\text{sub.2}}^2]$$

(78) To calculate the overall loss, the network combines the SAR loss and the amplitude loss, incorporating a weighting factor, α , for the amplitude loss. The total loss is computed as:

$$L_{\text{sub.total}} = L_{\text{sub.SAR}} + \alpha \times L_{\text{sub.amp}}$$

(79) The weighting factor value may be selected based on the dataset used during network training. In an embodiment, the network may be trained using two different SAR datasets: the National Geospatial-Intelligence Agency (NGA) SAR dataset and the Sandia National Laboratories Mini SAR Complex Imagery dataset, both of which feature complex-valued SAR images. In an embodiment, the weighting factor is set to 0.0001 for the NGA dataset and 0.00005 for the Sandia dataset. By integrating both the SAR and amplitude losses in the total loss function, the system effectively guides the training process to simultaneously address the removal of the artifacts and maintain the fidelity of the amplitude information. The weighting factor, α , enables AI deblocking network to balance the importance of the SAR loss and the amplitude loss, ensuring comprehensive optimization of the network during the training stages. In some implementations, diverse data augmentation techniques may be used to enhance the variety of training data. For example, techniques such as horizontal and vertical flops and rotations may be implemented on the training dataset. In an embodiment, model optimization is performed using MSE loss and Adam optimizer with a learning rate initially set to 1×10^{-4} and decreased by a factor of 2 at epochs 100, 200, and 250, with a total of 300 epochs. In an implementation, the batch size is set to 256×256 with each batch containing 16 images.

(80) Both branches first pass through a pixel unshuffling layer **211**, **221** which implements a pixel unshuffling process on the input data. Pixel unshuffling is a process used in image processing to reconstruct a high-resolution image from a low-resolution image by rearranging or “unshuffling” the pixels. The process can involve the following steps, low-resolution input, pixel arrangement, interpolation, and enhancement. The input to the pixel unshuffling algorithm is a low-resolution image (i.e., decompressed, quantized SAR I/Q data). This image is typically obtained by downscaling a higher-resolution image such as during the encoding process executed by encoder **110**. Pixel unshuffling aims to estimate the original high-resolution pixel values by redistributing and interpolating the low-resolution pixel values. The unshuffling process may involve performing interpolation techniques, such as nearest-neighbor, bilinear, or more sophisticated methods like bicubic or Lanczos interpolation, to estimate the missing pixel values and generate a higher-resolution image.

(81) The output of the unshuffling layers **211**, **221** may be fed into a series of layers which can include one or more convolutional layers and one or more parametric rectified linear unit (PReLU) layers. A legend is depicted for both FIG. 2A and FIG. 2B which indicates the cross hatched block represents a convolutional layer and the dashed block represents a PReLU layer. Convolution is the

first layer to extract features from an input image. Convolution preserves the relationship between pixels by learning image features using small squares of input data. It is a mathematical operation that takes two inputs such as an image matrix and a filter or kernel. The embodiment features a cascaded ResNet-like structure comprising 8 ResBlocks to effectively process the input data. The filter size associated with each convolutional layer may be different. The filter size used for the pixel domain of the top branch may be different than the filter size used for the frequency domain of the bottom branch.

(82) A PReLU layer is an activation function used in neural networks. The PReLU activation function extends the ReLU by introducing a parameter that allows the slope for negative values to be learned during training. The advantage of PReLU over ReLU is that it enables the network to capture more complex patterns and relationships in the data. By allowing a small negative slope for the negative inputs, the PReLU can learn to handle cases where the output should not be zero for all negative values, as is the case with the standard ReLU. In other implementations, other non-linear functions such as tanh or sigmoid can be used instead of PReLU.

(83) After passing through a series of convolutional and PReLU layers, both branches enter the resnet **230** which further comprises more convolutional and PReLU layers. The frequency domain branch is slightly different than the pixel domain branch once inside ResNet **230**, specifically the frequency domain is processed by a transposed convolutional (TConv) layer **231**. Transposed convolutions are a type of operation used in neural networks for tasks like image generation, image segmentation, and upsampling. They are used to increase the spatial resolution of feature maps while maintaining the learned relationships between features. Transposed convolutions aim to increase spatial dimensions of feature maps, effectively “upsampling” them. This is typically done by inserting zeros (or other values) between existing values to create more space for new values.

(84) Inside ResBlock **230** the data associated with the pixel and frequency domains are combined back into a single stream by using the output of the Tconv **231** and the output of the top branch. The combined data may be used as input for a channel-wise transformer **300**. In some embodiments, the channel-wise transformer may be implemented as a multi-scale attention block utilizing the attention mechanism. For more detailed information about the architecture and functionality of channel-wise transformer **300** refer to FIG. **3**. The output of channel-wise transformer **300** may be a bit stream suitable for reconstructing the original SAR I/Q image. FIG. **2B** shows the output of ResBlock **230** is passed through a final convolutional layer before being processed by a pixel shuffle layer **240** which can perform upsampling on the data prior to image reconstruction. The output of the AI deblocking network may be passed through a quantizer **124** for dequantization prior to producing a reconstructed SAR I/Q image **250**.

(85) FIG. **3** is a block diagram illustrating an exemplary architecture for a component of the system for SAR image compression, the channel-wise transformer **300**. According to the embodiment, channel-wise transformer receives an input signal, X_{in} **301**, the input signal comprising SAR I/Q data which is being processed by AI deblocking network **123**. The input signal may be copied and follow two paths through multi-channel transformer **300**.

(86) A first path may process input data through a position embedding module **330** comprising series of convolutional layers as well as a Gaussian Error Linear Unit (GeLU). In traditional recurrent neural networks or convolutional neural networks, the order of input elements is inherently encoded through the sequential or spatial nature of these architectures. However, in transformer-based models, where the attention mechanism allows for non-sequential relationships between tokens, the order of tokens needs to be explicitly conveyed to the model. Position embedding module **330** may represent a feedforward neural network (position-wise feedforward layers) configured to add position embeddings to the input data to convey the spatial location or arrangement of pixels in an image. The output of position embedding module **330** may be added to the output of the other processing path the received input signal is processed through.

(87) A second path may process the input data. It may first be processed via a channel-wise

configuration and then through a self-attention layer **320**. The signal may be copied/duplicated such that a copy of the received signal is passed through an average pool layer **310** which can perform a downsampling operation on the input signal. It may be used to reduce the spatial dimensions (e.g., width and height) of feature maps while retaining the most important information. Average pooling functions by dividing the input feature map into non-overlapping rectangular or square regions (often referred to as pooling windows or filters) and replacing each region with the average of the values within that region. This functions to downsample the input by summarizing the information within each pooling window.

(88) Self-attention layer **320** may be configured to provide an attention to AI deblocking network **123**. The self-attention mechanism, also known as intra-attention or scaled dot-product attention, is a fundamental building block used in various deep learning models, particularly in transformer-based models. It plays a crucial role in capturing contextual relationships between different elements in a sequence or set of data, making it highly effective for tasks involving sequential or structured data like complex-valued SAR I/Q channels. Self-attention layer **320** allows each element in the input sequence to consider other elements and weigh their importance based on their relevance to the current element. This enables the model to capture dependencies between elements regardless of their positional distance, which is a limitation in traditional sequential models like RNNs and LSTMs.

(89) The input **301** and downsampled input sequence is transformed into three different representations: Query (Q), Key (K), and Value (V). These transformations ($w_{sup.V}$, $w_{sup.K}$, and $w_{sup.Q}$) are typically linear projections of the original input. For each element in the sequence, the dot product between its Query and the Keys of all other elements is computed. The dot products are scaled by a factor to control the magnitude of the attention scores. The resulting scores may be normalized using a softmax function to get attention weights that represent the importance of each element to the current element. The Values (V) of all elements are combined using the attention weights as coefficients. This produces a weighted sum, where elements with higher attention weights contribute more to the final representation of the current element. The weighted sum is the output of the self-attention mechanism for the current element. This output captures contextual information from the entire input sequence.

(90) The output of the two paths (i.e., position embedding module **330** and self-attention layer **320**) may be combined into a single output data stream $x_{sub.out}$ **302**.

(91) FIG. 4 is a block diagram illustrating an exemplary system architecture **400** for providing lossless data compaction, according to an embodiment. As incoming data **401** is received by data deconstruction engine **402**. Data deconstruction engine **402** breaks the incoming data into sourceblocks, which are then sent to library manager **403**. Using the information contained in sourceblock library lookup table **404** and sourceblock library storage **405**, library manager **403** returns reference codes to data deconstruction engine **402** for processing into codewords, which are stored in codeword storage **106**. When a data retrieval request **407** is received, data reconstruction engine **408** obtains the codewords associated with the data from codeword storage **406**, and sends them to library manager **403**. Library manager **403** returns the appropriate sourceblocks to data reconstruction engine **408**, which assembles them into the proper order and sends out the data in its original form **409**.

(92) FIG. 5 is a diagram showing an embodiment of one aspect **500** of the system, specifically data deconstruction engine **501**. Incoming data **502** is received by data analyzer **503**, which optimally analyzes the data based on machine learning algorithms and input **504** from a sourceblock size optimizer, which is disclosed below. Data analyzer may optionally have access to a sourceblock cache **505** of recently processed sourceblocks, which can increase the speed of the system by avoiding processing in library manager **403**. Based on information from data analyzer **503**, the data is broken into sourceblocks by sourceblock creator **506**, which sends sourceblocks **507** to library manager **403** for additional processing. Data deconstruction engine **501** receives reference codes

508 from library manager **403**, corresponding to the sourceblocks in the library that match the sourceblocks sent by sourceblock creator **506**, and codeword creator **509** processes the reference codes into codewords comprising a reference code to a sourceblock and a location of that sourceblock within the data set. The original data may be discarded, and the codewords representing the data are sent out to storage **510**.

(93) FIG. **6** is a diagram showing an embodiment of another aspect of system **600**, specifically data reconstruction engine **601**. When a data retrieval request **602** is received by data request receiver **603** (in the form of a plurality of codewords corresponding to a desired final data set), it passes the information to data retriever **604**, which obtains the requested data **605** from storage. Data retriever **604** sends, for each codeword received, a reference codes from the codeword **606** to library manager **403** for retrieval of the specific sourceblock associated with the reference code. Data assembler **608** receives the sourceblock **607** from library manager **403** and, after receiving a plurality of sourceblocks corresponding to a plurality of codewords, assembles them into the proper order based on the location information contained in each codeword (recall each codeword comprises a sourceblock reference code and a location identifier that specifies where in the resulting data set the specific sourceblock should be restored to. The requested data is then sent to user **609** in its original form.

(94) FIG. **7** is a diagram showing an embodiment of another aspect of the system **700**, specifically library manager **701**. One function of library manager **701** is to generate reference codes from sourceblocks received from data deconstruction engine **701**. As sourceblocks are received **702** from data deconstruction engine **501**, sourceblock lookup engine **703** checks sourceblock library lookup table **704** to determine whether those sourceblocks already exist in sourceblock library storage **705**. If a particular sourceblock exists in sourceblock library storage **105**, reference code return engine **705** sends the appropriate reference code **706** to data deconstruction engine **601**. If the sourceblock does not exist in sourceblock library storage **105**, optimized reference code generator **407** generates a new, optimized reference code based on machine learning algorithms. Optimized reference code generator **707** then saves the reference code **708** to sourceblock library lookup table **704**; saves the associated sourceblock **709** to sourceblock library storage **105**; and passes the reference code to reference code return engine **705** for sending **706** to data deconstruction engine **501**. Another function of library manager **701** is to optimize the size of sourceblocks in the system. Based on information **711** contained in sourceblock library lookup table **404**, sourceblock size optimizer **410** dynamically adjusts the size of sourceblocks in the system based on machine learning algorithms and outputs that information **712** to data analyzer **603**. Another function of library manager **701** is to return sourceblocks associated with reference codes received from data reconstruction engine **601**. As reference codes are received **714** from data reconstruction engine **601**, reference code lookup engine **713** checks sourceblock library lookup table **715** to identify the associated sourceblocks; passes that information to sourceblock retriever **716**, which obtains the sourceblocks **717** from sourceblock library storage **405**; and passes them **718** to data reconstruction engine **601**.

(95) Detailed Description of Exemplary Aspects

(96) FIG. **8** is a flow diagram illustrating an exemplary method **800** for complex-valued SAR image compression, according to an embodiment. According to the embodiment, the process begins at step **801** when encoder **110** receives a raw complex-valued SAR image. The complex-valued SAR image comprises both I and Q components. In some embodiments, the I and Q components may be processed as separate channels. At step **802**, the received SAR image may be preprocessed for further processing by encoder **110**. For example, the input image may be clipped or otherwise transformed in order to facilitate further processing. As a next step **803**, the preprocessed data may be passed to quantizer **112** which quantizes the data. The next step **804**, comprises compressing the quantized SAR data using a compression algorithm known to those with skill in the art. In an embodiment, the compression algorithm may comprise HEVC encoding for both compression and decompression of SAR data. As a last step **805**, the compressed data may be compacted. The

compaction may be a lossless compaction technique, such as those described with reference to FIGS. 4-7. The output of method **800** is a compressed, compacted bit stream of SAR image data which can be stored in a database, requiring much less storage space than would be required to store the original, raw SAR image. The compressed and compacted bit stream may be transmitted to an endpoint for storage or processing. Transmission of the compressed and compacted data require less bandwidth and computing resources than transmitting raw SAR image data.

(97) FIG. **9** is a flow diagram illustrating an exemplary method **900** for decompression of a complex-valued SAR image, according to an embodiment. According to the embodiment, the process begins at step **901** when decoder **120** receives a bit stream comprising compressed and compacted complex-valued SAR image data. The compressed bit stream may be received from encoder **110** or from a suitable data storage device. At step **902**, the received bit stream is first de-compacted to produce an encoded (compressed) bit stream. In some embodiments, data reconstruction engine **601** may be implemented as a system for de-compacting a received bit stream. The next step **903**, comprising decompressing the de-compacted bit stream using a suitable compression algorithm known to those with skill in the art, such as HEVC encoding. At step **904**, the de-compressed SAR data may be fed as input into AI deblocking network **123** for image enhancement via a trained deep learning network. The AI deblocking network may utilize a series of convolutional layers and/or ResBlocks to process the input data and perform artifact removal on the de-compressed SAR image data. AI deblocking network may be further configured to implement an attention mechanism for the model to capture dependencies between elements regardless of their positional distance. In an embodiment, during training of AI deblocking network, the amplitude loss in conjunction with the SAR loss may be computed and accounted for, further boosting the compression performance of system **100**. The output of AI deblocking network **123** can be sent to a quantizer **124** which can execute step **905** by de-quantizing the output bit stream from AI deblocking network. As a last step **906**, system can reconstruct the original complex-valued SAR image using the de-quantized bit stream.

(98) FIG. **10** is a flow diagram illustrating an exemplary method for deblocking using a trained deep learning algorithm, according to an embodiment. According to the embodiment, the process begins at step **1001** wherein the trained deep learning algorithm (i.e., AI deblocking network **123**) receives a decompressed bit stream comprising SAR I/Q image data. At step **1002**, the bit stream is split into a pixel domain and a frequency domain. Each domain may pass through AI deblocking network, but have separate, almost similar processing paths. As a next step **1003**, each domain is processed through its respective branch, the branch comprising a series of convolutional layers and ResBlocks. In some implementations, frequency domain may be further processed by a transpose convolution layer. The two branches are combined and used as input for a multi-channel transformer with attention mechanism at step **1004**. Multi-channel transformer **300** may perform functions such as downsampling, positional embedding, and various transformations, according to some embodiments. Multi-channel transformer **300** may comprise one or more of the following components: channel-wise attention, transformer self-attention, and/or feedforward layers. In an implementation, the downsampling may be performed via average pooling. As a next step **1005**, the AI deblocking network processes the output of the channel-wise transformer. The processing may include the steps of passing the output through one or more convolutional or PRELU layers and/or upsampling the output. As a last step **1006**, the processed output may be forwarded to quantizer **124** or some other endpoint for storage or further processing.

(99) FIGS. **11A** and **11B** illustrate an exemplary architecture for an AI deblocking network configured to provide deblocking for a general N-channel data stream, according to an embodiment. The term “N-channel” refers to data that is composed of multiple distinct channels of modalities, where each channel represents a different aspect of type of information. These channels can exist in various forms, such as sensor readings, image color channels, or data streams, and they are often used together to provide a more comprehensive understanding of the underlying

phenomenon. Examples of N-channel data include, but is not limited to, RGB images (e.g., in digital images, the red, green, and blue channels represent different color information; combining these channels allows for the representation of a wide range of colors), medical imaging (e.g., may include Magnetic Resonance Imaging scans with multiple channels representing different tissue properties, or Computed Tomography scans with channels for various types of X-ray attenuation), audio data (e.g., stereo or multi-channel audio recordings where each channel corresponds to a different microphone or audio source), radar and lidar (e.g., in autonomous vehicles, radar and lidar sensors provide multi-channel data, with each channel capturing information about objects' positions, distances, and reflectivity) SAR image data, text data (e.g., in natural language processing, N-channel data might involve multiple sources of text, such as social media posts and news articles, each treated as a separate channel to capture different textual contexts), sensor networks (e.g., environmental monitoring systems often employ sensor networks with multiple sensors measuring various parameters like temperature, humidity, air quality, and more. Each sensor represents a channel), climate data, financial data, and social network data.

(100) The disclosed AI deblocking network may be trained to process any type of N-channel data, if the N-channel data has a degree of correlation. More correlation between and among the multiple channels yields a more robust and accurate AI deblocking network capable of performing high quality compression artifact removal on the N-channel data stream. A high degree of correlation implies a strong relationship between channels. Using SAR image data has been used herein as an exemplary use case for an AI deblocking network for a N-channel data stream comprising 2 channels, the In-phase and Quadrature components (i.e., I and Q, respectively).

(101) Exemplary data correlations that can be exploited in various implementations of AI deblocking network can include, but are not limited to, spatial correlation, temporal correlation, cross-sectional correlation (e.g., This occurs when different variables measured at the same point in time are related to each other), longitudinal correlation, categorical correlation, rank correlation, time-space correlation, functional correlation, and frequency domain correlation, to name a few.

(102) As shown, an N-channel AI deblocking network may comprise a plurality of branches **1110a-n**. The number of branches is determined by the number of channels associated with the data stream. Each branch may initially be processed by a series of convolutional and PRELU layers. Each branch may be processed by resnet **1130** wherein each branch is combined back into a single data stream before being input to N-channel wise transformer **1135**, which may be a specific configuration of transformer **300**. The output of N-channel wise transformer **1135** may be sent through a final convolutional layer before passing through a last pixel shuffle layer **1140**. The output of AI deblocking network for N-channel video/image data is the reconstructed N-channel data **1150**.

(103) As an exemplary use case, video/image data may be processed as a 3-channel data stream comprising Green (G), Red (R), and Blue (B) channels. An AI deblocking network may be trained that provides compression artifact removal of video/image data. Such a network would comprise 3 branches, wherein each branch is configured to process one of the three channels (R, G, or B). For example, branch **1110a** may correspond to the R-channel, branch **1110b** to the G-channel, and branch **1110c** to the B-channel. Each of these channels may be processed separately via their respective branches before being combined back together inside resnet **1130** prior to being processed by N-channel wise transformer **1135**.

(104) As another exemplary use case, a sensor network comprising a half dozen sensors may be processed as a 6-channel data stream. The exemplary sensor network may include various types of sensors collecting different types of, but still correlated, data. For example, sensor network can include a pressure sensor, a thermal sensor, a barometer, a wind speed sensor, a humidity sensor, and an air quality sensor. These sensors may be correlated to one another in at least one way. For example, the six sensors in the sensor network may be correlated both temporally and spatially, wherein each sensor provides a time series data stream which can be processed by one of the 6

channels **1110a-n** of AI deblocking network. As long as AI deblocking network is trained on N-channel data with a high degree of correlation and which is representative of the N-channel data it will encounter during model deployment, it can reconstruct the original data using the methods described herein.

(105) FIG. **12** is a block diagram illustrating an exemplary system architecture **1200** for N-channel data compression with predictive recovery, according to an embodiment. According to the embodiment, the system **1200** comprises an encoder module **1210** configured to receive as input N-channel data **1201** and compress and compact the input data into a bitstream **1202**, and a decoder module **1220** configured to receive and decompress the bitstream **1202** to output a reconstructed N-channel data **1203**.

(106) A data processor module **1211** may be present and configured to apply one or more data processing techniques to the raw input data to prepare the data for further processing by encoder **1210**. Data processing techniques can include (but are not limited to) any one or more of data cleaning, data transformation, encoding, dimensionality reduction, data slitting, and/or the like.

(107) After data processing, a quantizer **1212** performs uniform quantization on the n-number of channels. Quantization is a process used in various fields, including signal processing, data compression, and digital image processing, to represent continuous or analog data using a discrete set of values. It involves mapping a range of values to a smaller set of discrete values. Quantization is commonly employed to reduce the storage requirements or computational complexity of digital data while maintaining an acceptable level of fidelity or accuracy. Compressor **1213** may be configured to perform data compression on quantized N-channel data using a suitable conventional compression algorithm.

(108) The resulting encoded bitstream may then be (optionally) input into a lossless compactor (not shown) which can apply data compaction techniques on the received encoded bitstream. An exemplary lossless data compaction system which may be integrated in an embodiment of system **1200** is illustrated with reference to FIG. **4-7**. For example, lossless compactor may utilize an embodiment of data deconstruction engine **501** and library manager **403** to perform data compaction on the encoded bitstream. The output of the compactor is a compacted bitstream **1202** which can be stored in a database, requiring much less space than would have been necessary to store the raw N-channel data, or it can be transmitted to some other endpoint.

(109) At the endpoint which receives the transmitted compacted bitstream **1202** may be decoder module **1220** configured to restore the compacted data into the original SAR image by essentially reversing the process conducted at encoder module **1210**. The received bitstream may first be (optionally) passed through a lossless compactor which de-compacts the data into an encoded bitstream. In an embodiment, a data reconstruction engine **601** may be implemented to restore the compacted bitstream into its encoded format. The encoded bitstream may flow from compactor to decompressor **1222** wherein a data compaction technique may be used to decompress the encoded bitstream into the I/Q channels. It should be appreciated that lossless compactor components are optional components of the system, and may or may not be present in the system, dependent upon the embodiment.

(110) According to the embodiment, an Artificial Intelligence (AI) deblocking network **1223** is present and configured to utilize a trained deep learning network to provide compression artifact removal as part of the decoding process. AI deblocking network **1223** may leverage the relationship demonstrated between the various N-channels of a data stream to enhance the reconstructed N-channel data **1203**. Effectively, AI deblocking network **1223** provides an improved and novel method for removing compression artifacts that occur during lossy compression/decompression using a network designed during the training process to simultaneously address the removal of artifacts and maintain fidelity of the original N-channel data signal, ensuring a comprehensive optimization of the network during the training stages.

(111) The output of AI deblocking network **1223** may be dequantized by quantizer **1224**, restoring

the n-channels to their initial dynamic range. The dequantized n-channel data may be reconstructed and output **1203** by decoder module **1220** or stored in a database.

(112) FIG. **13** is a flow diagram illustrating an exemplary method for processing a compressed n-channel bit stream using an AI deblocking network, according to an embodiment. According to the embodiment, the process begins at step **1301** when a decoder module **1220** receives, retrieves, or otherwise obtains a bit stream comprising n-channel data with a high degree of correlation. At step **1302**, the bit stream is split into an n-number of domains. For example, if the received bit stream comprises image data in the form of R-, G-, and B-channels, then the bit stream would be split into 3 domains, one for each color (RGB). At step **1303**, each domain is processed through a branch comprising a series of convolutional layers and ResBlocks. The number of layers and composition of said layers may depend upon the embodiment and the n-channel data being processed. At step **1304**, the output of each branch is combined back into a single bitstream and used as an input into an n-channel wise transformer **1135**. At step **1305**, the output of the channel-wise transformer may be processed through one or more convolutional layers and/or transformation layers, according to various implementations. At step **1306**, the processed output may be sent to a quantizer for upscaling and other data processing tasks. As a last step **1307**, the bit stream may be reconstructed into its original uncompressed form.

(113) FIG. **14** is a block diagram illustrating a system for training a neural network to perform upsampling of decompressed data after lossy compression, according to an embodiment. The neural network may be referred to herein as a neural upsampler. According to the embodiment, a neural upsampler **1430** may be trained by taking training data **1402** which may comprise sets of two or more correlated datasets **101a-n** and performing whatever processing that is done to compress the data. This processing is dependent upon the type of data and may be different in various embodiments of the disclosed system and methods. For example, in the SAR imagery use case, the processing and lossy compression steps used quantization and HEVC compression of the I and Q images. The sets of compressed data may be used as input training data **1402** into the neural network **1420** wherein the target output is the original uncompressed data. Because there is correlation between the two or more datasets, the neural upsampler learns how to restore “lost” data by leveraging the cross-correlations.

(114) For each type of input data, there may be different compression techniques used, and different data conditioning for feeding into the neural upsampler. For example, if the input datasets **101a-n** comprise a half dozen correlated time series from six sensors arranged on a machine, then delta encoding or a swinging door algorithm may be implemented for data compression and processing.

(115) The neural network **1420** may process the training data **1402** to generate model training output in the form of restored dataset **1430**. The neural network output may be compared against the original dataset to check the model's precision and performance. If the model output does not satisfy a given criteria or some performance threshold, then parametric optimization **1415** may occur wherein the training parameters and/or network hyperparameters may be updated and applied to the next round of neural network training.

(116) FIG. **15** is a flow diagram illustrating an exemplary method **1500** for training a neural network to perform upsampling of decompressed data after lossy compression, according to an embodiment. According to an embodiment, the process begins at step **1501** by creating a training dataset comprising compressed data by performing lossy compression on two or more datasets which are substantially correlated. As a next step **1502**, the training dataset is used to train a neural network (i.e., neural upsampler) configured to leverage the correlation between the two or more datasets to generate as output a reconstructed dataset. At step **1503**, the output of the neural network is compared to the original two more datasets to determine if the performance of the neural network at reconstructing the compressed data. If the model performance is not satisfactory, which may be determined by a set of criteria or some performance metric or threshold, then the neural

network model parameters and/or hyperparameters may be updated **1504** and applied to the next round of training as the process moves to step **1502** and iterates through the method again.

(117) FIG. **16** is a block diagram illustrating an exemplary architecture for a neural upsampler configured to process N-channel time-series data, according to an embodiment. The neural upsampler may comprise a trained deep learning algorithm. According to the embodiment, a neural upsampler configured to process time-series data may comprise a recurrent autoencoder with an n-channel transformer attention network. In such an embodiment, the neural upsampler may be trained to process decompressed time-series data wherein the output of the upsampler is restored time-series data (i.e., restore most of the lost data due to the lossy compression). The upsampler may receive decompressed n-channel time-series data comprising two or more data sets of time-series data which are substantially correlated. For example, the two or more data sets may comprise multiple sets of Internet of Things (IoT) sensor data from sensors that are likely to be temporally correlated. For instance, consider a large number of sensors on a single complex machine (e.g., a combine tractor, a 3D printer, construction equipment, etc.) or a large number of sensors in a complex systems such as a pipeline or refinery.

(118) The n-channel time-series data may be received split into separate channels **1610a-n** to be processed individually by encoder **1620**. In some embodiments, encoder **1620** may employ a series of various data processing layers which may comprise recurrent neural network (RNN) layers, pooling layers, PRELU layers, and/or the like. In some implementations, one or more of the RNN layers may comprise a Long Short-Term Memory (LSTM) network. In some implementations, one or more of the RNN layers may comprise a sequence-to-sequence model. In yet another implementation, the one or more RNN layer may comprise a gate recurrent unit (GRU). Each channel may be processed by its own series of network layers wherein the encoder **1620** can learn a representation of the input data which can be used to determine the defining features of the input data. Each individual channel then feeds into an n-channel wise transformer **1630** which can learn the interdependencies between the two or more channels of correlated time-series data. The output of the n-channel wise transformer **1630** is fed into the decoder **1640** component of the recurrent autoencoder in order to restore missing data lost due to a lossy compression implemented on the time-series data. N-channel wise transformer **1630** is designed so that it can weigh the importance of different parts of the input data and then capture long-range dependencies between and among the input data. The decoder may process the output of the n-channel wise transformer **1630** into separate channels comprising various layers as described above. The output of decoder **1640** is the restored time-series data **1602**, wherein most of the data which was “lost” during lossy compression can be recovered using the neural upsampler which leverages the interdependencies hidden within correlated datasets.

(119) In addition to RNNs and their variants, other neural network architectures like CNNs and hybrid models that combine CNNs and RNNs can also be implemented for processing time series and sensor data, particularly when dealing with sensor data that can be structured as images or spectrograms. For example, if you had, say, 128 time series streams, it could be structured as two 64×64 pixel images (64 times series each, each with 64 time steps), and then use the same approach as the described above with respect to the SAR image use case. In an embodiment, a one-dimensional CNN can be used as a data processing layer in encoder **1620** and/or decoder **1640**. The selection of the neural network architecture for time series data processing may be based on various factors including, but not limited to, the length of the input sequences, the frequency and regularity of the data points, the need to handle multivariate input data, the presence of exogenous variables or covariates, the computational resources available, and/or the like.

(120) The exemplary time-series neural upsampler described in FIG. **16** may be trained on a training dataset comprising a plurality of compressed time-series data sourced from two or more datasets which are substantially correlated. For example, in a use case directed towards neural upsampling of IoT sensor data, the neural upsampler may be trained on a dataset comprising

compressed IoT sensor data. During training, the output of the neural upsampler may be compared against the non-compressed version of the IoT sensor data to determine the neural upsampler's performance on restoring lost information.

(121) FIG. **17** is a block diagram illustrating an exemplary system architecture **1700** for upsampling of decompressed sensor data after lossy compression using a neural network, according to an embodiment. According to the embodiment, a neural upsampler **1730** is present and configured to receive decompressed sensor data (e.g., time-series data obtained from an IoT device) and restore the decompressed data by leveraging learned data correlations and inter- and intra-dependencies. According to an embodiment, the system may receive a plurality of sensor data **1701a-n** from two or more sensors/devices, wherein the sensor data are substantially correlated. In an embodiment, the plurality of sensor data **1701a-n** comprises time-series data. Time-series data received from two or more sensors may be temporally correlated, for example, IoT data from a personal fitness device and a blood glucose monitoring device during the time when a user of both devices is exercising may be correlated in time and by heart rate. As another example, a large number of sensors used to monitor a manufacturing facility may be correlated temporally.

(122) A data compressor **1710** is present and configured to utilize one or more data compression methods on received sensor data **1701a-n**. The data compression method chosen must be a lossy compression method. Exemplary types of lossy compression that may be used in some embodiments may be directed towards image or audio compression such as JPEG and MP3, respectively. For time series data lossy compression methods that may be implemented include (but is not limited to) one or more of the following: delta encoding, swinging door algorithm, batching, data aggregation, feature extraction. In an implementation, data compressor **1710** may implement network protocols specific for IoT such as message queuing telemetry transport (MQTT) for supporting message compression on the application layer and/or constrained application protocol (CoAP) which supports constrained nodes and networks and can be used with compression.

(123) The compressed multi-channel sensor data **1701a-n** may be decompressed by a data decompressor **1720** which can utilize one or more data decompression methods known to those with skill in the art. The output of data decompressor **1720** is a sensor data stream(s) of decompressed data which is missing information due to the lossy nature of the compression/decompression methods used. The decompressed sensor data stream(s) may be passed to neural upsampler **1730** which can utilize a trained neural network to restore most of the "lost" information associated with the decompressed sensor data stream(s) by leveraging the learned correlation(s) between and among the various sensor data streams. The output of neural upsampler **1730** is restored sensor data **1740**.

(124) FIG. **18** is a flow diagram illustrating an exemplary method **1800** for performing neural upsampling of two or more time-series data streams, according to an embodiment. In this example, the two or more time-series streams may be associated with large sets of IoT sensors/devices. The two or more time-series streams are substantially correlated. The two or more time-series data streams may be temporally correlated. For example, a plurality of IoT sensors may be time-synchronized to better understand cause-and-effect relationships.

(125) A neural upsampler which has been trained on compressed time-series data associated with one or more IoT sensor channels is present and configured to restore time-series data which has undergone lossy data compression and decompression by leveraging the correlation between the sensor data streams. A non-exhaustive list of time-series data correlations that may be used by an embodiment of the system and method can include cross-correlation and auto-correlation.

(126) The two or more time-series data streams may be processed by a data compressor **1710** employing a lossy compression method. The lossy compression method may implement a lossy compression algorithm appropriate for compressing time-series data. The choice of compression implementation may be based on various factors including, but not limited to, the type of data being processed, the computational resources and time required, and the use case of the upsampler.

Exemplary time-series data compression techniques which may be used include, but are not limited to, delta encoding, swinging door algorithm data aggregation, feature extraction, and batching, to name a few. The compressed time series data may be store in a database and/or transmitted to an endpoint. The compressed time-series data may be sent to a data decompressor **1720** which may employ a lossy decompression technique on the compressed time-series data. The decompressed data may be sent to the neural upsampler which can restore the decompressed data to nearly its original state by leveraging the temporal (and/or other) correlation between the time-series IoT sensor data streams. The compressed time-series data is received by data decompressor **1720** at step **1801**. At data decompressor **1720** the compressed time-series data may be decompressed via a lossy decompression algorithm at step **1802**.

(127) A neural upsampler for restoration of time-series (e.g., IoT sensor data) data received from two or more data channels may be trained using two or more datasets comprising compressed time-series data which is substantially correlated. For example, the two or more datasets may comprise time-series data from a plurality of sensors affixed to a long-haul semi-truck and configured to monitor various aspects of the vehicles operation and maintenance and report the monitored data to a central data processing unit which can compress and transmit the data for storage or further processing. The two or more sensor channels are correlated in various ways such as temporally. In various embodiments, each channel of the received time-series data may be fed into its own neural network comprising a series of convolutional and/or recurrent and ReLU and/or pooling layers which can be used to learn latent correlations in the feature space that can be used to restore data which has undergone lossy compression. A multi-channel transformer may be configured to receive the output of each of the neural networks produce, learn from the latent correlation in the feature space, and produce reconstructed time-series data. At step **1803**, the decompressed time-series data may be used as input to the trained neural upsampler configured to restore the lost information of the decompressed time-series data. The neural upsampler can process the decompressed data to generate as output restored time-series data at step **1804**.

(128) FIG. **19** is a block diagram illustrating an exemplary system architecture for neural upsampling of two or more genomic datasets, according to an embodiment. Genomic data **1910a-n** may comprise, for example, any one or more of DNA sequences, single nucleotide polymorphisms (SNPs) gene expression data, epigenetic data, structural genomic data, mitochondrial DNA sequences, and/or the like. These examples highlight different layers of genomic information, from the basic DNA sequence to variations, gene expression, and epigenetic modifications. Analyzing and integrating multiple types of genomic data are crucial for a comprehensive understanding of biological processes, evolution, and the genetic basis of diseases. Thus, it would be beneficial to have a system, method, and/or computer readable instructions capable of providing neural upsampling of genomic data (e.g., human genomes or subsets of them, any parallel genome data sets, two or more persons mitochondrial DNA sequences, etc.) which has undergone lossy compression, therefore nearly restoring all the lost data.

(129) In an embodiment, genomic data **1910a-n** may comprise parallel genome datasets. Parallel genome datasets typically refer to multiple sets of genomic data that are generated or analyzed simultaneously. Using parallel sequencing runs, multiple samples may undergo DNA sequencing simultaneously in parallel, generating multiple sets of sequencing data concurrently. For example, in a genomics laboratory, several DNA samples might be processed and sequenced using high-throughput sequencing technologies in a single sequencing run, producing parallel datasets. In another example, genomic data from different individuals or populations may be collected and analyzed concurrently to study genetic diversity, population structure, and evolutionary patterns. Researchers might analyze genome sequences from individuals of different ethnicities or geographic regions in parallel to investigate population-specific genetic variations.

(130) There are several common data formats used for storing and transmitting genomic data, and which may be used in various implementations of the disclosed system and methods. These formats

are designed to efficiently represent the vast amount of information generated through various genomic technologies. One such format of genomic data which may be processed by system **1900** is Format for Sequence Data (FASTA). FASTA is a text-based format for representing nucleotide or protein sequences. It consists of a header line starting with “>”, followed by the sequence data. This format may be used when processing genomic data such as DNA, RNA, and protein sequences. Similarly, Format for Quality Scores (FASTQ) may be used in some implementations. FASTQ is a text-based format that extends FASTA by including quality scores for each base in the sequence. It is commonly used for storing data from next-generation sequencing (NGS) platforms. (131) Another exemplary format which may be processed by system **1900** is sequence alignment/mapping (SAM/BAM). SAM is a text-based format for representing sequence alignment data, while BAM is the binary equivalent. SAM/BAM files store aligned sequencing reads along with quality scores, mapping positions, and other relevant information. SAM/BAM may be implemented in use cases for storing and exchanging data related to sequence alignments, such as is the case in the context of NGS data. As a final example, variant call format (VCF) may be implemented in some embodiments of system **1900**. VCF is a text-based format for representing genomic variations, such as single nucleotide polymorphisms (SNPs), insertions, deletions, and structural variants.

(132) The genomic data may be received at a data compressor **1920** which is present and configured to utilize one or more data compression methods on received genomic data **1910a-n**. Genomic data, especially raw sequencing data, can be massive, and compression techniques are often employed to reduce storage requirements and facilitate data transfer. The data compression method chosen must be a lossy compression method. Exemplary types of lossy compression that may be used in some embodiments include quality score quantization, reference-based compression, subsampling, genomic data transformation, and lossy compression of read data. (133) In an embodiment where quality score quantization is implemented, quality scores associated with each base in sequencing data represent the confidence in the accuracy of the base call. These scores are often encoded with high precision, but for compression purposes, they can be quantized to reduce the bit depth, introducing a level of information loss. Higher quantization levels reduce the precision of quality scores but can significantly reduce file sizes.

(134) In an embodiment, where reference-based compression is implemented, instead of storing the entire genomic sequence, some compression methods store only the differences between the target sequence and a reference genome. Variations and mutations are encoded, while the reference genome provides a framework. This method can achieve substantial compression, but some specific information about the individual's genome is lost. Raw read data from sequencing platforms may contain redundant or noisy information. Lossy compression algorithms may filter or smooth the data to reduce redundancy or noise. While this can result in higher compression, it may lead to the loss of some information, especially in regions with lower sequencing quality.

(135) Genomic data compressed by data compressor **1920** may then be sent to a data decompressor **1930** which can utilize one or more data decompression methods known to those with skill in the art. The output of data decompressor **1930** is a genomic data stream(s) of decompressed data which is missing information due to the lossy nature of the compression/decompression methods used. The decompressed genomic data stream(s) may be passed to neural upsampler **1940** which can utilize a trained neural network to restore most of the “lost” information associated with the decompressed genomic data stream(s) by leveraging the learned correlation(s) between and among the various genomic datasets. The output of neural upsampler **1940** is restored genomic data **1950**.

(136) According to various embodiments, system **1900** utilizes a trained neural upsampler to leverage correlations in the received two or more genomic datasets **1910a-n** in order to restore lost data. In an implementations, neural upsampler **1940** may comprise a series of recurrent neural network layers, pooling layers, an n-channel transformer, and/or convolutional layers as described herein. In an embodiment, neural upsampler **1940** may be trained on a training dataset comprising a

corpus of compressed genomic data, wherein the compressed genomic data is correlated. The neural upsampler may be trained to generate as output genomic data, which is at close to its original state, prior to undergoing lossy data compression. The genomic data which was used to create the training dataset may be kept and used to validate the training output of neural upsampler, in this way the neural upsampler can be trained to generate output which nearly matches the original, uncompressed genomic data.

(137) Genomic datasets can be correlated with each other in various ways, providing valuable insights into biological relationships, evolutionary history, and disease associations. There are some ways in which distinct genomic datasets can be correlated, and which may be learned and leveraged by a trained neural upsampler **1940** to restore genomic data which has been processed via lossy compression/decompression. For example, genetic variation and linkage disequilibrium can provide correlation between and among genetic datasets **1910a-n**. SNPs are variations at a single nucleotide position in the DNA sequence. Correlating SNP data across different genomic datasets can reveal patterns of genetic variation and linkage disequilibrium. Haplotype blocks found in genomic data may be used as a learned correlation by neural upsampler. Haplotypes are combinations of alleles on a single chromosome. Understanding the correlation of haplotypes across datasets helps in identifying linked genetic variations. Yet another correlation that can be found among genetic datasets is phenotypic correlation. Correlating genomic data with phenotypic information can identify genetic variants associated with specific traits or diseases. This is commonly done through Genome-Wide Association Studies (GWAS) and can involve comparing different genomic datasets.

(138) More examples of genetic data correlations which may be leveraged in one or more embodiments include evolutionary relationships, gene expression correlation, epigenetic correlations, structural genomic correlation, functional annotations, and population genetics. Human mitochondrial DNA (mtDNA) sequences can be correlated to one another in several ways to understand genetic relationships, population structure, and evolutionary history. Some common approaches for analyzing and correlating human mitochondrial sequences can include phylogenetic analysis, haplogroup assignment, and population genetics and diversity measures. Phylogenetic trees are constructed based on sequence differences, revealing the evolutionary relationships among different mitochondrial haplotypes. This is often done using methods like Maximum Likelihood or Bayesian inference. Phylogenetic trees help identify clades, lineages, and common ancestors, providing insights into the historical relationships among mitochondrial sequences. Mitochondrial DNA is categorized into haplogroups, which represent major branches of the mitochondrial phylogenetic tree. Haplogroups are defined by specific polymorphisms and sequence variations. Assigning individuals to haplogroups allows for broader categorization of mtDNA diversity and helps trace maternal lineages. A neural upsampler can use the correlations in genomic datasets to be trained to restore lost data.

(139) FIG. **20** is a flow diagram illustrating an exemplary method **2000** for performing neural upsampling of two or more genomic datasets, according to an embodiment. In this example, the two or more genomic datasets (also referred to as data streams) may be associated with human genomic data (e.g., human genome). The two or more genomic datasets are substantially correlated as described herein. For example, two or more people's mitochondrial DNA sequences will be closely related.

(140) A neural upsampler which has been trained on compressed genomic data is present and configured to restore time-series data which has undergone lossy data compression and decompression by leveraging the correlation between the genomic datasets. A non-exhaustive list of genomic data correlations that may be used by an embodiment of the system and method can include genetic variation and linkage disequilibrium, and haplotype blocks.

(141) The two or more genomic datasets may be processed by a data compressor **1920** employing a lossy compression method. The lossy compression method may implement a lossy compression

algorithm appropriate for compressing genomic data. The choice of compression implementation may be based on various factors including, but not limited to, the type of data being processed, the computational resources and time required, and the use case of the upsampler. Exemplary genomic data compression techniques which may be used include, but are not limited to, quality score quantization, reference-based compression, subsampling, and genomic data transformation, to name a few. The compressed genomic data may be stored in a database and/or transmitted to an endpoint. The compressed genomic data may be sent to a data decompressor **1930** which may employ a lossy decompression technique on the compressed genomic data. The decompressed data may be sent to the neural upsampler which can restore the decompressed data to nearly its original state by leveraging the genetic variation (and/or other) correlation between the genomic datasets. The compressed genomic data is received by data decompressor **1930** at step **2001**. At data decompressor **1930** the compressed genomic data may be decompressed via a lossy decompression algorithm at step **2002**.

(142) A neural upsampler for restoration of genomic (e.g., human genomes or subsets thereof) data received from two or more data channels may be trained using two or more datasets comprising compressed genomic data which is substantially correlated. For example, the two or more datasets may comprise genomic data from a subset of the human genome. Subsets of human genomes refer to specific groups or categories of genetic information within the larger human population. These subsets can be defined based on various criteria, such as geographical origin, shared genetic features, or clinical characteristics. Here are some examples of subsets of human genomes: haplogroups, population specific genomic variation, ancestral populations, ethnic and geographical groups, disease-specific subsets, founder populations (i.e., groups of individuals who established a new population, often with a limited gene pool), isolate populations, age-specific subsets, long-lived individuals, and/or the like. The two or more subsets of human genomes are correlated in various ways such as temporally. In various embodiments, each channel of the received genomic data may be fed into its own neural network comprising a series of convolutional and/or recurrent and ReLU and/or pooling layers which can be used to learn latent correlations in the feature space that can be used to restore data which has undergone lossy compression. A multi-channel transformer may be configured to receive the output that each of the neural networks produce, learn from the latent correlation in the feature space, and produce reconstructed genomic data. At step **2003**, the decompressed genomic data may be used as input to the trained neural upsampler configured to restore the lost information of the decompressed genomic data. The neural upsampler can process the decompressed data to generate as output restored genomic data at step **2004**.

(143) FIG. **22** is a block diagram illustrating an exemplary architecture for ProteomicUNet, a neural network configured to upsample compressed proteomic data, according to an embodiment. ProteomicUNet combines convolutional layers, LSTM layers, and attention mechanisms to effectively restore high-resolution proteomic data from compressed inputs. The architecture is designed to leverage both spatial and temporal dependencies in proteomic data, making it particularly well-suited for handling complex proteomic signals.

(144) According to an embodiment, ProteomicUNet receives compressed proteomic data **2201** as input. This compressed data may be in the form of a reduced-dimension matrix or a sparse representation, typically obtained through lossy compression techniques such as dimensionality reduction or quantization applied to original high-resolution proteomic data. The input data first passes through a series of convolutional layers **2220**, which extract local patterns and features from the compressed data. These convolutional layers apply learned filters to capture spatial dependencies and create hierarchical representations of the proteomic signals. Multiple convolutional layers are stacked, with increasing numbers of filters and decreasing spatial dimensions, to capture multi-scale features.

(145) The output of the convolutional layers is then processed by one or more LSTM layers **2230**. These LSTM layers are crucial for capturing temporal dependencies and long-range interactions in

the proteomic data. LSTMs are particularly effective at modeling the complex dynamics and relationships between different proteomic channels or time points, preserving contextual information and capturing the overall structure of the proteomic signals.

(146) An attention mechanism **2240** is incorporated into the architecture to focus on the most informative regions and features of the compressed proteomic data. This mechanism learns to assign different weights to various parts of the input data based on their relevance and importance for the upsampling task. By prioritizing the most salient features, the attention mechanism improves the quality of the reconstructed high-resolution data.

(147) The outputs from the LSTM layers and the attention mechanism are then fed into a series of upsampling layers **2250**. These layers gradually increase the spatial resolution of the proteomic data, employing techniques such as transposed convolutions, bilinear interpolation, or pixel shuffling. The upsampling layers learn to reconstruct high-resolution proteomic signals by combining the learned features and attention weights from the previous layers.

(148) ProteomicUNet incorporates skip connections **2260** to facilitate information flow between the downsampling (convolutional) and upsampling paths. These connections allow the model to combine high-level semantic information from the downsampling path with localized details from the upsampling path, helping to preserve fine-grained structures and reduce information loss during the compression and upsampling process.

(149) The final output of ProteomicUNet is a reconstructed high-resolution proteomic dataset **2251** that closely resembles the original uncompressed data. The model is trained to minimize the difference between this reconstructed data and the ground truth using appropriate loss functions, such as mean squared error or cross-entropy loss.

(150) FIG. **23** is a method diagram illustrating an exemplary method for training ProteomicUNet, according to an embodiment. The process begins at step **2301** with the preparation of a training dataset comprising paired compressed and uncompressed proteomic data. This dataset is crucial for supervised learning, allowing the model to learn the mapping from compressed representations to their high-resolution counterparts.

(151) At step **2302**, the ProteomicUNet model is initialized with random weights. The model architecture is set up as described in the previous paragraphs, including the convolutional layers, LSTM layers, attention mechanism, upsampling layers, and skip connections.

(152) The training process begins at step **2303**, where the model processes batches of compressed proteomic data and generates upsampled predictions. These predictions are compared to the corresponding uncompressed ground truth data using appropriate loss functions. The model parameters are then optimized using gradient descent algorithms such as Adam or RMSprop, with appropriate learning rate schedules.

(153) Step **2304** involves the application of data augmentation techniques to enhance the model's robustness and generalization ability. These techniques may include random cropping, flipping, or noise injection applied to the training data.

(154) The trained model's performance is evaluated at step **2305** using various metrics such as peak signal-to-noise ratio (PSNR), structural similarity index (SSIM), or cosine similarity. These metrics assess the quality of the reconstructed proteomic data. Cross-validation techniques are employed to assess the model's generalization ability and prevent overfitting.

(155) At step **2306**, the trained ProteomicUNet is utilized to upsample new compressed proteomic data. The compressed data is input into the model, processed through the various layers, and a high-resolution reconstruction is generated as output.

(156) FIG. **24** is a method diagram illustrating an exemplary method for utilizing the trained ProteomicUNet, according to an embodiment. The process begins at step **2401** with the input of compressed proteomic data that requires upsampling.

(157) At step **2402**, the compressed proteomic data is processed through the convolutional layers of the ProteomicUNet. These layers apply learned filters to capture spatial dependencies and create

hierarchical representations of the proteomic signals. Multiple convolutional layers are used, with increasing numbers of filters and decreasing spatial dimensions, to capture multi-scale features. (158) In step **2403**, the output from the convolutional layers is fed into one or more LSTM layers. These LSTM layers capture temporal dependencies and long-range interactions in the proteomic data, modeling the complex dynamics and relationships between different proteomic channels or time points.

(159) Step **2404** involves the application of the attention mechanism. This mechanism assigns different weights to various parts of the input data based on their relevance and importance for the upsampling task, focusing on the most informative regions and features of the compressed proteomic data.

(160) At step **2405**, the data passes through a series of upsampling layers. These layers gradually increase the spatial resolution of the proteomic data, employing techniques such as transposed convolutions or bilinear interpolation. The upsampling layers reconstruct high-resolution proteomic signals by combining the learned features and attention weights from the previous layers.

(161) In step **2406**, the skip connections are utilized. These connections combine high-level semantic information from the earlier layers with localized details from the upsampling path, helping to preserve fine-grained structures and reduce information loss during the upsampling process.

(162) Finally, at step **2407**, the model outputs the reconstructed high-resolution proteomic dataset. This output closely resembles the original uncompressed data and can be used for further analysis, visualization, or integration with other data types in the broader genomic and proteomic analysis pipeline.

(163) FIG. **25** is a block diagram illustrating an exemplary architecture for ProteinFormer, a novel protein structure prediction model based on the transformer architecture, according to an embodiment. ProteinFormer is designed to handle protein sequences using amino acid embeddings and positional encodings, leveraging the power of self-attention mechanisms to capture long-range dependencies and structural patterns in protein sequences.

(164) According to the embodiment, ProteinFormer begins with an input representation subsystem **2510**. This subsystem receives protein sequences **2501** and tokenizes them into amino acid tokens using a vocabulary of 20 standard amino acids plus special tokens for rare amino acids and padding. Each amino acid token is then converted into a learnable embedding vector of dimension **512** that captures its biochemical properties such as hydrophobicity, charge, and size. Sinusoidal positional encodings are added to these embeddings to preserve the sequential order of amino acids in the protein sequence. These encodings are calculated using sine and cosine functions of different frequencies for each dimension of the embedding vector.

(165) The embedded and position-encoded input is then processed by the transformer encoder **2520**, which forms the core of the ProteinFormer architecture. The transformer encoder consists of 6 identical layers, each containing two sub-layers: multi-head self-attention **2521** and feed-forward neural networks **2522**. The multi-head self-attention mechanism uses 8 attention heads, each with a dimensionality of 64, allowing each amino acid to attend to other amino acids in the sequence, effectively capturing long-range dependencies and structural relationships. The feed-forward networks in each layer consist of two linear transformations with a ReLU activation in between, with an inner-layer dimensionality of **2048**. Layer normalization is applied before each sub-layer, and residual connections are used around each of the two sub-layers.

(166) Following the transformer encoder, the structure prediction subsystem **2530** processes the encoder's output to predict the 3D coordinates of each amino acid in the protein structure. This subsystem **2530** employs a combination of distance prediction and torsion angle prediction. First, it predicts pairwise distances between Ca atoms using a 2D convolutional neural network applied to the outer product of the encoder's output. Then, it predicts phi and psi torsion angles for each residue using a bidirectional LSTM network. These predicted distances and angles are then used as

constraints in a differentiable geometry optimization procedure to generate the final 3D coordinates.

(167) ProteinFormer incorporates a novel self-supervised pre-training task, where the model predicts masked amino acids in a protein sequence. In this task, 15% of the input amino acids are randomly masked, replaced with a special [MASK] token. The model is then trained to predict the original amino acids at these masked positions. This pre-training approach allows the model to learn long-range dependencies and structural patterns from large-scale protein sequence databases, even in the absence of experimentally determined structures. The pre-training is performed on a dataset of millions of protein sequences from diverse organisms.

(168) The output of ProteinFormer is a predicted 3D protein structure **2531**, represented as a series of (x, y, z) coordinates for each Ca atom in the protein backbone. Additionally, the model outputs confidence scores for each predicted position, allowing downstream tasks to account for prediction uncertainty. The predicted structure is provided in a standard PDB (Protein Data Bank) format, which can be directly used in various molecular visualization tools and for further analysis in structural biology and drug discovery pipelines.

(169) FIG. **26** is a method diagram illustrating an exemplary method for utilizing ProteinFormer for protein structure prediction, according to an embodiment. The process begins at step **2601** with the input of a protein sequence. This sequence is typically represented as a string of amino acid letters, corresponding to the primary structure of the protein.

(170) At step **2602**, the input protein sequence is preprocessed. This involves tokenizing the sequence into individual amino acid tokens using a vocabulary of 20 standard amino acids plus special tokens for rare amino acids and padding. Each token is then converted into its corresponding embedding vector of dimension **512**, capturing biochemical properties such as hydrophobicity, charge, and size. Sinusoidal positional encodings are added to these embeddings to preserve the sequential information of the amino acids in the protein.

(171) In step **2603**, the embedded and position-encoded sequence is passed through the 6 transformer encoder layers of ProteinFormer. Each layer applies 8-head self-attention mechanisms and feed-forward neural networks with inner-layer dimensionality of **2048**, allowing the model to capture complex relationships and dependencies within the protein sequence. Layer normalization and residual connections are applied in each layer.

(172) Following the encoder, step **2604** involves the structure prediction subsystem. This subsystem takes the output of the transformer encoder and generates predictions for the 3D coordinates of each amino acid in the protein structure. It first predicts pairwise distances between Ca atoms using a 2D convolutional neural network, then predicts phi and psi torsion angles for each residue using a bidirectional LSTM network.

(173) At step **2605**, the predicted distances and angles are used as constraints in a differentiable geometry optimization procedure to assemble the complete protein structure model. This step may involve additional refinement processes to ensure the physical feasibility of the predicted structure, such as energy minimization or clash resolution.

(174) The final step **2606** involves outputting the predicted 3D protein structure. This structure is represented as a series of (x, y, z) coordinates for each Ca atom in the protein backbone, along with confidence scores for each predicted position. The structure is provided in standard PDB (Protein Data Bank) format, allowing for easy visualization and further analysis in structural biology and drug discovery pipelines.

(175) FIG. **27** is a block diagram illustrating an exemplary architecture for BindingSiteFinder, a graph neural network (GNN) based model for predicting protein binding sites, according to an embodiment. BindingSiteFinder integrates information from upsampled proteomic data, predicted 3D structures, and evolutionary conservation to accurately identify potential binding sites on proteins.

(176) According to an embodiment, BindingSiteFinder begins with an input feature subsystem

2710. This subsystem processes three key inputs: protein sequences **2701**, predicted 3D structures **2702**, and evolutionary conservation scores **2703**. Protein sequences **2701** are encoded using learnable amino acid embeddings. The predicted 3D structures **2702**, typically obtained from models such as ProteinFormer, are represented as graphs where nodes correspond to amino acids and edges capture spatial proximity and physicochemical interactions. Evolutionary conservation scores **2703**, derived from multiple sequence alignments, are incorporated as additional node features in the graph representation.

(177) The core of BindingSiteFinder is the Graph Neural Network (GNN) subsystem **2720**. This subsystem processes the protein structure graphs to learn representations that capture both local and global structural patterns relevant to binding site prediction. The GNN layers propagate information between neighboring nodes **2721**, allowing the model to consider the spatial arrangement and interdependencies of amino acids. Attention mechanisms **2722** may be employed within the GNN to focus on the most informative regions and interactions within the protein structure.

(178) Following the GNN subsystem, the binding site classification subsystem **2730** processes the learned graph representations. This subsystem consists of a series of fully connected neural network layers that classify each amino acid as either part of a binding site or not. The first layer takes the GNN subsystem's output as input, with subsequent layers gradually reducing dimensionality. The architecture employs dropout for regularization and ReLU activation functions in intermediate layers. The final layer outputs either binary classifications or probability scores for each amino acid, using sigmoid or softmax activation depending on the desired output format. The classification can be performed using binary cross-entropy loss or focal loss to handle potential class imbalance, as binding sites typically comprise a small portion of the total protein surface.

(179) BindingSiteFinder incorporates a specialized message-passing scheme **2740** that efficiently integrates features from the upsampled proteomic data and evolutionary conservation scores. This scheme facilitates information flow across the protein structure, enabling the model to capture complex patterns that may indicate potential binding sites. The message-passing operation updates each node's features by aggregating information from neighboring nodes and edges. It collects and transforms feature vectors from the neighborhood, aggregates them using functions like sum or mean, combines this with the node's own features, and applies a non-linear activation. This process iterates multiple times, allowing comprehensive information propagation throughout the protein structure and capturing both local and global patterns relevant to binding site prediction.

(180) The output of BindingSiteFinder is a prediction of binding sites **2741** on the input protein structure. These predictions can be represented as a binary mask over the protein structure or as probability scores for each amino acid, indicating the likelihood of it being part of a binding site.

(181) FIG. **28** is a flow diagram illustrating an exemplary method for utilizing BindingSiteFinder to predict protein binding sites, according to an embodiment. The process begins at step **2801** with the input of three elements: a protein sequence, a predicted 3D protein structure (typically obtained from models like ProteinFormer), and evolutionary conservation scores derived from multiple sequence alignments. The protein structure is typically represented as a 3D coordinate system of atoms, often in a standard format such as PDB (Protein Data Bank).

(182) At step **2802**, these inputs are preprocessed and converted into a graph representation. In this graph, nodes represent amino acids and edges represent spatial proximity and interactions between amino acids based on the predicted 3D structure. The protein sequence is encoded using learnable amino acid embeddings. The evolutionary conservation scores are incorporated as additional node features in the graph representation. Additional features, such as physicochemical properties of amino acids, are also included as node attributes.

(183) In step **2803**, the graph representation is processed through the GNN layers of BindingSiteFinder. Each layer in the GNN applies message passing operations, allowing information to propagate between neighboring nodes and capturing both local and global structural

patterns relevant to binding site prediction.

(184) Following the GNN layers, step **2804** involves the binding site classification subsystem. This subsystem takes the learned graph representations and applies fully connected layers to classify each amino acid as either part of a binding site or not. The classification process considers both the structural information captured by the GNN and the additional features incorporated in the graph representation.

(185) At step **2805**, the model's predictions are post-processed. This may involve applying a threshold to the classification scores, clustering predicted binding site residues, or refining the predictions based on known structural or chemical constraints.

(186) The final step **2806** involves outputting the predicted binding sites. These can be represented in various formats, such as a list of amino acid residues likely to be involved in binding, a 3D visualization of the protein structure with highlighted binding sites, or a set of probability scores for each residue.

(187) FIG. **29** is a block diagram illustrating an exemplary integrated system architecture combining ProteinFormer and BindingSiteFinder for comprehensive protein structure and binding site prediction, according to an embodiment. This integrated system creates a powerful pipeline that leverages the strengths of both models to provide a more complete analysis of protein structures and their functional sites.

(188) According to the embodiment, the integrated system begins with the ProteinFormer subsystem **2500**. This subsystem takes a protein sequence as input and generates a predicted 3D structure using its transformer-based architecture. The predicted structure includes the spatial coordinates of each amino acid in the protein, representing its tertiary structure.

(189) The output from ProteinFormer is then seamlessly passed to the BindingSiteFinder subsystem **2700**. This subsystem takes the predicted 3D structure as one of its primary inputs, along with the original protein sequence and evolutionary conservation data. The evolutionary conservation data may be accessible through a database **2904**, in an embodiment. The predicted structure is converted into a graph representation, with nodes representing amino acids and edges representing spatial proximity and interactions.

(190) A data integration layer **2910** ensures compatibility between the outputs of ProteinFormer and the input requirements of BindingSiteFinder. This layer may perform tasks such as format conversion, normalization of coordinate systems, and integration of additional features like evolutionary conservation scores that are required by BindingSiteFinder but not produced by ProteinFormer.

(191) The BindingSiteFinder subsystem then processes the integrated data through its graph neural network architecture (GNN subsystem **2710**) to predict potential binding sites on the protein structure. The binding site classification subsystem **2720** then processes these predictions. The output of this subsystem is a set of predicted binding sites, typically represented as probability scores for each amino acid or as clusters of amino acids likely to form binding sites.

(192) A post-processing and visualization subsystem **2920** takes the outputs from both ProteinFormer and BindingSiteFinder to create comprehensive visualizations and reports. This subsystem may generate 3D visualizations of the predicted protein structure with highlighted binding sites, produce detailed reports on the structural features and predicted functional sites, and provide confidence scores for both the overall structure prediction and the binding site predictions.

(193) The integrated system also includes a feedback loop **2930** that allows the predictions from BindingSiteFinder to inform and potentially improve the structure predictions of ProteinFormer in subsequent iterations. This feedback mechanism might adjust the structure prediction to better accommodate the predicted binding sites, leading to more biologically relevant structural models.

(194) The final output of the integrated system **2921** is a comprehensive protein analysis package that includes the predicted 3D structure, identified potential binding sites, and associated confidence scores and metrics. This output can be used for various downstream applications in

structural biology, drug discovery, and protein engineering.

(195) The integration of these two sophisticated models presents several technical challenges that are addressed in the system design. These include ensuring consistent data representations between the two models, managing computational resources efficiently to handle the intensive processing required by both models, and developing robust error handling and uncertainty quantification methods that account for the compounded uncertainties from both structure prediction and binding site identification.

(196) By combining the strengths of ProteinFormer in predicting overall protein structure with BindingSiteFinder's ability to identify potential binding sites, this integrated system provides a powerful tool for researchers in the fields of structural biology and drug discovery. It enables a more complete understanding of protein structures and their functional sites from sequence information alone, potentially accelerating the pace of discovery in these fields.

(197) FIG. **30** is a method diagram illustrating an exemplary method for the integrated use of ProteinFormer and BindingSiteFinder for comprehensive protein structure and binding site prediction, according to an embodiment. This method outlines the sequential steps involved in processing a protein sequence through both models to obtain a final prediction of structure and binding sites.

(198) The process begins at step **3001** with the input of a protein sequence, typically represented as a string of amino acid letters corresponding to the primary structure of the protein. This input protein sequence is processed through ProteinFormer. This step involves tokenizing the sequence, applying embeddings and positional encodings, and then passing the data through the transformer encoder layers. The output of this step is a predicted 3D structure of the protein.

(199) In step **3002**, the predicted 3D structure from ProteinFormer is validated and, if necessary, refined. This may involve energy minimization, clash resolution, or other techniques to ensure the physical feasibility of the predicted structure.

(200) Step **3003** involves preparing the ProteinFormer output for input into BindingSiteFinder. This includes converting the 3D structure into a graph representation, where nodes represent amino acids and edges represent spatial relationships and interactions. Additional features, such as evolutionary conservation scores, are also incorporated at this stage.

(201) At step **3004**, the prepared data is processed through BindingSiteFinder. The graph neural network layers of BindingSiteFinder analyze the structural and evolutionary information to predict potential binding sites on the protein surface.

(202) Step **3005** involves post-processing the binding site predictions. This may include clustering predicted binding site residues, applying confidence thresholds, or refining predictions based on known biochemical constraints.

(203) In step **3006**, the results from both ProteinFormer and BindingSiteFinder are integrated and visualized. This step produces comprehensive 3D visualizations of the predicted protein structure with highlighted binding sites, as well as detailed reports on structural features and predicted functional sites.

(204) The final step **3007** involves outputting the comprehensive protein analysis results. This output includes the predicted 3D structure, identified potential binding sites, associated confidence scores, and any additional metrics or annotations generated during the analysis process.

(205) This integrated method leverages the strengths of both ProteinFormer and BindingSiteFinder, providing a more complete analysis of protein structure and function from sequence information alone. The seamless flow of information between these two sophisticated models enables researchers to gain deeper insights into protein structures and their potential interaction sites, facilitating advancements in fields such as structural biology and drug discovery.

(206) FIG. **31** is a block diagram illustrating the structure of ProteoGraph, a multi-modal graph representation designed to integrate various types of proteomic data, according to an embodiment. ProteoGraph serves as a central data structure that facilitates the analysis, interpretation, and

visualization of complex proteomic information.

(207) The system begins with input data **3101**, which includes protein sequences, mass spectrometry data, predicted 3D structures, functional annotations, and other relevant proteomic information.

(208) This input **3101** is first processed by a data integration and preprocessing subsystem **3105** that combines information from various sources, including proteomic experiments, protein databases, structural databases, and binding site databases. This subsystem applies quality control measures, normalization techniques, and data filtering to ensure the consistency and reliability of the integrated data. The preprocessed data is then used to construct the initial graph structure **3170**.

(209) According to an embodiment, ProteoGraph consists of several types of nodes, each representing different proteomic entities. Protein nodes **3110** represent individual proteins, containing attributes such as protein sequences, functional annotations, and experimental metadata. Peptide nodes **3120** represent peptides identified from mass spectrometry experiments and are connected to their corresponding protein nodes. Modification nodes **3130** represent post-translational modifications (PTMs) detected on peptides and are linked to their associated peptide nodes. Structure nodes **3140** represent predicted 3D structures of proteins, obtained from models like ProteinFormer, and are connected to their corresponding protein nodes. Binding site nodes **3150** represent predicted binding sites, identified by methods such as BindingSiteFinder, and are linked to their associated structure nodes.

(210) The relationships between these nodes are represented by edges in the ProteoGraph. Protein-peptide edges **3111** connect protein nodes to their corresponding peptide nodes. Peptide-modification edges **3107** link peptide nodes to their associated modification nodes. Protein-structure edges **3108** connect protein nodes to their predicted 3D structure nodes. Structure-binding site edges **3109** link structure nodes to their predicted binding site nodes. Protein-binding site edges **3110** directly connect protein nodes to their predicted binding site nodes, providing a direct link between proteins and their functional sites.

(211) Each node and edge in ProteoGraph carries specific attributes that provide detailed information about the proteomic entities and their relationships. For example, protein nodes may include attributes such as expression levels and functional annotations, while peptide nodes might contain mass-to-charge ratios and retention times. Edge attributes can include confidence scores for the relationships they represent or quantitative measurements such as peptide abundance.

(212) The constructed graph is then passed to the graph embedding and analysis subsystem **3160**, which applies advanced techniques such as Graph Convolutional Networks (GCNs) to learn low-dimensional representations of the nodes and edges in ProteoGraph. The GCN architecture used in ProteoGraph consists of multiple graph convolutional layers, each followed by a rectified linear unit (ReLU) activation function. The first layer takes as input the node feature matrix X and the adjacency matrix A of the graph. Each subsequent layer performs the following operation:

$$H(l+1) = \sigma(D^{-1/2} \hat{A} D^{-1/2} H(l) W(l))$$

where $H(l)$ is the matrix of activations in the l -th layer, $\hat{A} = A + I$ is the adjacency matrix with added self-connections, D is the degree matrix of $\hat{A}$, $W(l)$ is the weight matrix for layer l , and σ is the ReLU activation function.

(213) The output of ProteoGraph **3161** consists of multi-modal graph representations that integrate various types of proteomic data. These representations include the graph structure itself, with nodes representing proteins, peptides, modifications, structures, and binding sites, and edges representing their relationships. Each node and edge carries attributes providing detailed proteomic information. Additionally, ProteoGraph outputs the low-dimensional embeddings generated by the graph embedding and analysis subsystem **3160**, which capture the structural and functional relationships between the proteomic entities. These embeddings can be used for various downstream analysis tasks such as protein function prediction, interaction prediction, or clustering of similar proteins. The output can be provided in standard graph file formats such as GraphML or JSON for further

processing or visualization in external tools. The GCN is trained using a semi-supervised learning approach. A subset of nodes in the graph is labeled with known properties (e.g., protein functions, binding site locations), and the model is trained to predict these properties for the labeled nodes while simultaneously learning informative embeddings for all nodes. The loss function combines a supervised component (e.g., cross-entropy loss for labeled nodes) and an unsupervised component (e.g., a graph regularization term that encourages connected nodes to have similar embeddings).

(214) The training process uses mini-batch stochastic gradient descent with the Adam optimizer. To handle the large scale of the ProteoGraph, we employ a neighborhood sampling strategy during training, where for each target node, we sample a fixed number of its neighbors to form a computational subgraph.

(215) These embeddings can then be used for various downstream tasks such as protein function prediction, interaction prediction, or clustering of similar proteins based on their multi-modal representations.

(216) FIG. **32** is a method diagram illustrating an exemplary method for constructing and utilizing ProteoGraph, according to an embodiment. The process begins at step **3201** with the collection and integration of diverse proteomic data sources. These may include protein sequences, mass spectrometry data, predicted 3D structures, functional annotations, and other relevant proteomic information. The collected data undergoes preprocessing and quality control in the data integration and preprocessing subsystem. This involves normalizing data from different sources, filtering out low-quality or unreliable data points, and ensuring consistency across the dataset.

(217) In step **3202**, the preprocessed data is used to construct the ProteoGraph structure. This involves creating nodes for each proteomic entity (proteins, peptides, modifications, structures, and binding sites) and establishing edges to represent the relationships between these entities.

(218) Step **3203** involves the annotation of nodes and edges with relevant attributes. This includes adding detailed information such as protein sequences, peptide properties, modification types, structural coordinates, and confidence scores to the appropriate nodes and edges.

(219) At step **3203**, graph embedding techniques are applied to the constructed ProteoGraph. The graph embedding and analysis subsystem uses a Graph Convolutional Network (GCN) to learn low-dimensional representations of the nodes and edges, capturing the complex relationships within the proteomic data.

(220) Step **3204** involves the analysis and interpretation of the ProteoGraph. This may include tasks such as protein function prediction, identification of protein-protein interactions, or discovery of potential drug targets using the learned graph embeddings.

(221) The final step **3205** focuses on the output of the Proteograph subsystem. The output includes the final graph structure with all nodes and edges, their attributes, and the learned node embeddings.

(222) By representing proteomic data in this comprehensive, interconnected manner, ProteoGraph provides a powerful framework for integrating and analyzing complex proteomic information. It enables researchers to uncover hidden patterns, generate hypotheses, and gain deeper insights into protein structures, functions, and interactions.

(223) FIG. **33** is a block diagram illustrating the architecture of the deep learning system with latent transformer core, according to an embodiment. This system is designed to process large codeword models in genomic and proteomic data analysis efficiently while maintaining high fidelity in data representation.

(224) According to an embodiment, the system begins with an input layer **3301** that receives a plurality of input vectors. These vectors may represent genomic or proteomic data, such as DNA sequences, protein structures, or mass spectrometry data. The input data is first converted into sourceblocks, which are then assigned codewords based on a dictionary generated by a codebook generation subsystem.

(225) The input vectors are first processed by a variational autoencoder (VAE) encoder **3310**. This

encoder compresses the input data into a latent space representation, reducing dimensionality while preserving key features of the data. The VAE encoder captures essential features and variations of the input data, creating a compact and informative latent space vector. The heart of the system is the latent transformer **3320**, which operates on the latent space vectors produced by the VAE encoder. Unlike traditional transformers, this component does not include embedding or positional encoding layers, as these aspects are inherently captured in the latent space representation. The latent transformer applies self-attention mechanisms to learn relationships between different elements of the latent vectors, capturing complex dependencies in the data.

(226) The output of the latent transformer is then processed by the VAE decoder **3330**, which reconstructs the data from the latent space back into the original data space **3331**. This reconstruction process helps ensure that the learned representations maintain fidelity to the original data. The VAE decoder can either reconstruct the original input data or generate new data based on the processed latent vectors.

(227) The deep learning system with latent transformer core may serve as an intermediary step in the genomic and proteomic data analysis pipeline. It interfaces with the ProteomicUNet **2200**, ProteinFormer **2500**, and BindingSiteFinder **2700** components to enhance their performance and enable more efficient processing of large-scale data. Following the initial upsampling and reconstruction of compressed proteomic data by ProteomicUNet **2200**, the latent transformer system processes this high-dimensional output. It encodes the data into a compact latent space representation using its variational autoencoder, then applies transformer-based processing to capture complex relationships within this latent space.

(228) This approach offers several advantages: it reduces dimensionality for more efficient downstream processing, extracts relevant features that can enhance the performance of subsequent analyses, and enables effective handling of large-scale proteomic datasets. The resulting processed latent vectors serve as optimized input for ProteinFormer **2500** and BindingSiteFinder **2700**, potentially improving the accuracy of 3D structure predictions and binding site identifications.

(229) Additionally, this architecture allows for the integration of genomic data alongside proteomic data in a unified latent space representation, facilitating more comprehensive, multi-omic analyses. By positioning the latent transformer system at this juncture in the pipeline, we enhance the overall scalability and effectiveness of our proteomic analysis framework, particularly when dealing with complex, large-scale datasets. After processing by ProteinFormer **2500** and BindingSiteFinder **2700**, the results can be fed back through the latent transformer system, converting them back to the original data space for final output and visualization in the ProteoGraph **3100** representation.

(230) The entire system can be trained end-to-end, optimizing the VAE encoder, latent transformer, and VAE decoder simultaneously. It uses techniques like gradient descent and backpropagation to minimize reconstruction loss and capture relevant patterns in the latent space. This allows the system to handle various tasks, including data compression, denoising, anomaly detection, and data generation across multiple domains within genomic and proteomic data analysis.

(231) FIG. **34** is a method diagram illustrating the use of latent space transformation for genomic and proteomic data. The process begins at step **3401** with the input of genomic or proteomic data. This could include DNA sequences, protein structures, or mass spectrometry data. At step **3402**, the input data is encoded into a latent space representation using the VAE encoder. This step compresses the data while preserving its essential features.

(232) In step **3403**, the latent space representations undergo transformation through the latent transformer. This step captures complex relationships and dependencies within the data. Step **3404** involves decoding the transformed latent representations back into the original data space using the VAE decoder. Finally, at step **3505**, the system outputs the processed genomic or proteomic data, which now incorporates the learned relationships and transformations.

(233) This process enables the deep learning system with latent transformer core to efficiently handle large-scale genomic and proteomic datasets, capturing complex patterns and dependencies

while maintaining high fidelity in data representation.

(234) FIG. **35** is a block diagram illustrating data flow through components of the system, in an embodiment. The diagram begins with raw genomic and proteomic data input **3501**. This data first passes through the ProteomicUNet **2200** for upsampling and enhancement. The enhanced proteomic data, along with genomic sequences, are then fed into the deep learning system with latent transformer core **3300**. This system processes the data through its variational autoencoder and transformer components, creating optimized latent space representations.

(235) The output from the deep learning system **3300** is then passed to ProteinFormer **2500** for 3D structure prediction. The predicted structures, along with the enhanced proteomic data, are then processed by BindingSiteFinder **2700** to identify potential binding sites.

(236) All this information—the enhanced proteomic data, latent space representations, predicted structures, and identified binding sites—is integrated into the ProteoGraph **3100** representation.

(237) By integrating these sophisticated components—ProteomicUNet, the deep learning system with latent transformer core, ProteinFormer, BindingSiteFinder, and ProteoGraph—this system provides a powerful and comprehensive framework for advanced genomic and proteomic data analysis. It enables researchers to extract deeper insights from complex biological data, potentially accelerating discoveries in fields such as structural biology, functional genomics, and drug discovery.

(238) The system and methods described herein have a wide range of potential applications in genomics, proteomics, and integrated multi-omics analysis. As a non-limiting example, the ProteomicUNet component may be employed in biomarker discovery research. Consider a scenario where a research team is analyzing low-resolution mass spectrometry data from blood samples of patients with a rare autoimmune disease. The data quality might be insufficient for reliable biomarker identification using traditional methods. By applying ProteomicUNet to upsample the compressed, low-resolution mass spectrometry data, the team could potentially generate high-resolution proteomic profiles that reveal previously undetectable peptide peaks and post-translational modifications. This application of ProteomicUNet could allow researchers to extract more information from existing data, potentially identifying novel biomarkers without requiring additional expensive experiments.

(239) In another non-limiting example, the ProteinFormer component could be utilized in rapidly evolving public health scenarios, such as predicting structures of new SARS-COV-2 variants. Researchers needing to quickly predict the structures of spike proteins from emerging variants to assess their potential impact on vaccine efficacy could input the amino acid sequences of new spike protein variants into ProteinFormer. The system could provide accurate 3D structural predictions of the variant spike proteins, highlighting structural changes that could affect antibody binding. This application of ProteinFormer could offer rapid structural insights without the need for time-consuming experimental structure determination, allowing for faster assessment of new variants in critical situations.

(240) The BindingSiteFinder component, as another non-limiting example, could be applied in pharmaceutical research for identifying new binding sites for drug development. A pharmaceutical company searching for new drug targets to treat Alzheimer's disease might have a list of proteins associated with the disease but lack information on potential binding sites. By using BindingSiteFinder to analyze the structures of these proteins (which could be predicted by ProteinFormer), the company could generate a list of predicted binding sites on each protein, ranked by likelihood and including characteristics of each site. This application of

(241) BindingSiteFinder could accelerate the drug discovery process by providing potential binding sites for virtual screening, potentially reducing the time and cost of initial drug candidate identification.

(242) In a non-limiting example of ProteoGraph's application, it could be used in personalized medicine for integrating multi-omics data for cancer treatment. Oncologists with genomic,

proteomic, and clinical data from a cohort of cancer patients might struggle to integrate this information for personalized treatment decisions. By using ProteoGraph to create a comprehensive, multi-modal representation of each patient's data, the medical team could generate an integrated graph for each patient, visualizing relationships between genetic mutations, protein expression levels, predicted protein structures, and clinical outcomes. This use of ProteoGraph could allow for intuitive visualization and analysis of complex, multi-modal data, potentially facilitating personalized treatment decisions based on a holistic view of each patient's molecular profile.

(243) The deep learning system with latent transformer core, in a non-limiting example, could be applied to efficiently process large-scale genomic data. A national biobank with petabytes of genomic data from millions of individuals might need to efficiently process this data for population-wide studies while maintaining privacy. By using the deep learning system with latent transformer core to process the large-scale genomic data, the biobank could generate compressed representations of genomic data that preserve key features and relationships, potentially allowing for efficient storage and analysis. This system could enable the biobank to perform large-scale genomic analyses with reduced computational resources while maintaining data privacy through the use of latent space representations.

(244) As a non-limiting example of the integrated system's use, consider its application in the comprehensive analysis of a novel protein family. Biologists discovering a new family of proteins potentially involved in cellular aging, with limited experimental data but needing comprehensive insights into these proteins' structures, functions, and interactions, could employ the full integrated system. Starting with genomic and limited proteomic data for these proteins, the research team could use ProteomicUNet to upsample the limited proteomic data, ProteinFormer to predict 3D structures for the protein family, BindingSiteFinder to identify potential binding sites on these structures, and ProteoGraph to integrate all the data into a comprehensive representation. The deep learning system with latent transformer core could then process the integrated data to identify patterns and relationships. This integrated approach could provide a wealth of information and hypotheses about the novel protein family, potentially guiding future experimental work and accelerating the research process.

(245) These examples illustrate some of the potential applications of the system and its components, but it should be understood that the system's use is not limited to these scenarios and can be adapted to a wide range of research and clinical contexts in genomics, proteomics, and multi-omics integration.

(246) Exemplary Computing Environment

(247) FIG. 21 illustrates an exemplary computing environment on which an embodiment described herein may be implemented, in full or in part. This exemplary computing environment describes computer-related components and processes supporting enabling disclosure of computer-implemented embodiments. Inclusion in this exemplary computing environment of well-known processes and computer components, if any, is not a suggestion or admission that any embodiment is no more than an aggregation of such processes or components. Rather, implementation of an embodiment using processes and components described in this exemplary computing environment will involve programming or configuration of such processes and components resulting in a machine specially programmed or configured for such implementation. The exemplary computing environment described herein is only one example of such an environment and other configurations of the components and processes are possible, including other relationships between and among components, and/or absence of some processes or components described. Further, the exemplary computing environment described herein is not intended to suggest any limitation as to the scope of use or functionality of any embodiment implemented, in whole or in part, on components or processes described herein.

(248) The exemplary computing environment described herein comprises a computing device **10** (further comprising a system bus **11**, one or more processors **20**, a system memory **30**, one or more

interfaces **40**, one or more non-volatile data storage devices **50**), external peripherals and accessories **60**, external communication devices **70**, remote computing devices **80**, and cloud-based services **90**.

(249) System bus **11** couples the various system components, coordinating operation of and data transmission between, those various system components. System bus **11** represents one or more of any type or combination of types of wired or wireless bus structures including, but not limited to, memory busses or memory controllers, point-to-point connections, switching fabrics, peripheral busses, accelerated graphics ports, and local busses using any of a variety of bus architectures. By way of example, such architectures include, but are not limited to, Industry Standard Architecture (ISA) busses, Micro Channel Architecture (MCA) busses, Enhanced ISA (EISA) busses, Video Electronics Standards Association (VESA) local busses, a Peripheral Component Interconnects (PCI) busses also known as a Mezzanine busses, or any selection of, or combination of, such busses. Depending on the specific physical implementation, one or more of the processors **20**, system memory **30** and other components of the computing device **10** can be physically co-located or integrated into a single physical component, such as on a single chip. In such a case, some or all of system bus **11** can be electrical pathways within a single chip structure.

(250) Computing device may further comprise externally-accessible data input and storage devices **12** such as compact disc read-only memory (CD-ROM) drives, digital versatile discs (DVD), or other optical disc storage for reading and/or writing optical discs **62**; magnetic cassettes, magnetic tape, magnetic disk storage, or other magnetic storage devices; or any other medium which can be used to store the desired content and which can be accessed by the computing device **10**.

Computing device may further comprise externally-accessible data ports or connections **12** such as serial ports, parallel ports, universal serial bus (USB) ports, and infrared ports and/or transmitter/receivers. Computing device may further comprise hardware for wireless communication with external devices such as IEEE 1394 (“Firewire”) interfaces, IEEE 802.11 wireless interfaces, BLUETOOTH® wireless interfaces, and so forth. Such ports and interfaces may be used to connect any number of external peripherals and accessories **60** such as visual displays, monitors, and touch-sensitive screens **61**, USB solid state memory data storage drives (commonly known as “flash drives” or “thumb drives”) **63**, printers **64**, pointers and manipulators such as mice **65**, keyboards **66**, and other devices **67** such as joysticks and gaming pads, touchpads, additional displays and monitors, and external hard drives (whether solid state or disc-based), microphones, speakers, cameras, and optical scanners.

(251) Processors **20** are logic circuitry capable of receiving programming instructions and processing (or executing) those instructions to perform computer operations such as retrieving data, storing data, and performing mathematical calculations. Processors **20** are not limited by the materials from which they are formed, or the processing mechanisms employed therein, but are typically comprised of semiconductor materials into which many transistors are formed together into logic gates on a chip (i.e., an integrated circuit or IC). The term processor includes any device capable of receiving and processing instructions including, but not limited to, processors operating on the basis of quantum computing, optical computing, mechanical computing (e.g., using nanotechnology entities to transfer data), and so forth. Depending on configuration, computing device **10** may comprise more than one processor. For example, computing device **10** may comprise one or more central processing units (CPUs) **21**, each of which itself has multiple processors or multiple processing cores, each capable of independently or semi-independently processing programming instructions. Further, computing device **10** may comprise one or more specialized processors such as a graphics processing unit (GPU) **22** configured to accelerate processing of computer graphics and images via a large array of specialized processing cores arranged in parallel.

(252) System memory **30** is processor-accessible data storage in the form of volatile and/or nonvolatile memory. System memory **30** may be either or both of two types: non-volatile memory

and volatile memory. Non-volatile memory **30a** is not erased when power to the memory is removed, and includes memory types such as read only memory (ROM), electronically-erasable programmable memory (EEPROM), and rewritable solid state memory (commonly known as “flash memory”). Non-volatile memory **30a** is typically used for long-term storage of a basic input/output system (BIOS) **31**, containing the basic instructions, typically loaded during computer startup, for transfer of information between components within computing device, or a unified extensible firmware interface (UEFI), which is a modern replacement for BIOS that supports larger hard drives, faster boot times, more security features, and provides native support for graphics and mouse cursors. Non-volatile memory **30a** may also be used to store firmware comprising a complete operating system **35** and applications **36** for operating computer-controlled devices. The firmware approach is often used for purpose-specific computer-controlled devices such as appliances and Internet-of-Things (IoT) devices where processing power and data storage space is limited. Volatile memory **30b** is erased when power to the memory is removed and is typically used for short-term storage of data for processing. Volatile memory **30b** includes memory types such as random access memory (RAM), and is normally the primary operating memory into which the operating system **35**, applications **36**, program modules **37**, and application data **38** are loaded for execution by processors **20**. Volatile memory **30b** is generally faster than non-volatile memory **30a** due to its electrical characteristics and is directly accessible to processors **20** for processing of instructions and data storage and retrieval. Volatile memory **30b** may comprise one or more smaller cache memories which operate at a higher clock speed and are typically placed on the same IC as the processors to improve performance.

(253) Interfaces **40** may include, but are not limited to, storage media interfaces **41**, network interfaces **42**, display interfaces **43**, and input/output interfaces **44**. Storage media interface **41** provides the necessary hardware interface for loading data from non-volatile data storage devices **50** into system memory **30** and storage data from system memory **30** to non-volatile data storage device **50**. Network interface **42** provides the necessary hardware interface for computing device **10** to communicate with remote computing devices **80** and cloud-based services **90** via one or more external communication devices **70**. Display interface **43** allows for connection of displays **61**, monitors, touchscreens, and other visual input/output devices. Display interface **43** may include a graphics card for processing graphics-intensive calculations and for handling demanding display requirements. Typically, a graphics card includes a graphics processing unit (GPU) and video RAM (VRAM) to accelerate display of graphics. One or more input/output (I/O) interfaces **44** provide the necessary support for communications between computing device **10** and any external peripherals and accessories **60**. For wireless communications, the necessary radio-frequency hardware and firmware may be connected to I/O interface **44** or may be integrated into I/O interface **44**.

(254) Non-volatile data storage devices **50** are typically used for long-term storage of data. Data on non-volatile data storage devices **50** is not erased when power to the non-volatile data storage devices **50** is removed. Non-volatile data storage devices **50** may be implemented using any technology for non-volatile storage of content including, but not limited to, CD-ROM drives, digital versatile discs (DVD), or other optical disc storage; magnetic cassettes, magnetic tape, magnetic disc storage, or other magnetic storage devices; solid state memory technologies such as EEPROM or flash memory; or other memory technology or any other medium which can be used to store data without requiring power to retain the data after it is written. Non-volatile data storage devices **50** may be non-removable from computing device **10** as in the case of internal hard drives, removable from computing device **10** as in the case of external USB hard drives, or a combination thereof, but computing device will typically comprise one or more internal, non-removable hard drives using either magnetic disc or solid state memory technology. Non-volatile data storage devices **50** may store any type of data including, but not limited to, an operating system **51** for providing low-level and mid-level functionality of computing device **10**, applications **52** for providing high-level functionality of computing device **10**, program modules **53** such as

containerized programs or applications, or other modular content or modular programming, application data **54**, and databases **55** such as relational databases, non-relational databases, and graph databases.

(255) Applications (also known as computer software or software applications) are sets of programming instructions designed to perform specific tasks or provide specific functionality on a computer or other computing devices. Applications are typically written in high-level programming languages such as C++, Java, and Python, which are then either interpreted at runtime or compiled into low-level, binary, processor-executable instructions operable on processors **20**. Applications may be containerized so that they can be run on any computer hardware running any known operating system. Containerization of computer software is a method of packaging and deploying applications along with their operating system dependencies into self-contained, isolated units known as containers. Containers provide a lightweight and consistent runtime environment that allows applications to run reliably across different computing environments, such as development, testing, and production systems.

(256) The memories and non-volatile data storage devices described herein do not include communication media. Communication media are means of transmission of information such as modulated electromagnetic waves or modulated data signals configured to transmit, not store, information. By way of example, and not limitation, communication media includes wired communications such as sound signals transmitted to a speaker via a speaker wire, and wireless communications such as acoustic waves, radio frequency (RF) transmissions, infrared emissions, and other wireless media.

(257) External communication devices **70** are devices that facilitate communications between computing device and either remote computing devices **80**, or cloud-based services **90**, or both. External communication devices **70** include, but are not limited to, data modems **71** which facilitate data transmission between computing device and the Internet **75** via a common carrier such as a telephone company or internet service provider (ISP), routers **72** which facilitate data transmission between computing device and other devices, and switches **73** which provide direct data communications between devices on a network. Here, modem **71** is shown connecting computing device **10** to both remote computing devices **80** and cloud-based services **90** via the Internet **75**. While modem **71**, router **72**, and switch **73** are shown here as being connected to network interface **42**, many different network configurations using external communication devices **70** are possible. Using external communication devices **70**, networks may be configured as local area networks (LANs) for a single location, building, or campus, wide area networks (WANs) comprising data networks that extend over a larger geographical area, and virtual private networks (VPNs) which can be of any size but connect computers via encrypted communications over public networks such as the Internet **75**. As just one exemplary network configuration, network interface **42** may be connected to switch **73** which is connected to router **72** which is connected to modem **71** which provides access for computing device **10** to the Internet **75**. Further, any combination of wired **77** or wireless **76** communications between and among computing device **10**, external communication devices **70**, remote computing devices **80**, and cloud-based services **90** may be used. Remote computing devices **80**, for example, may communicate with computing device through a variety of communication channels **74** such as through switch **73** via a wired **77** connection, through router **72** via a wireless connection **76**, or through modem **71** via the Internet **75**. Furthermore, while not shown here, other hardware that is specifically designed for servers may be employed. For example, secure socket layer (SSL) acceleration cards can be used to offload SSL encryption computations, and transmission control protocol/internet protocol (TCP/IP) offload hardware and/or packet classifiers on network interfaces **42** may be installed and used at server devices.

(258) In a networked environment, certain components of computing device **10** may be fully or partially implemented on remote computing devices **80** or cloud-based services **90**. Data stored in non-volatile data storage device **50** may be received from, shared with, duplicated on, or offloaded

to a non-volatile data storage device on one or more remote computing devices **80** or in a cloud computing service **92**. Processing by processors **20** may be received from, shared with, duplicated on, or offloaded to processors of one or more remote computing devices **80** or in a distributed computing service **93**. By way of example, data may reside on a cloud computing service **92**, but may be usable or otherwise accessible for use by computing device **10**. Also, certain processing subtasks may be sent to a microservice **91** for processing with the result being transmitted to computing device **10** for incorporation into a larger processing task. Also, while components and processes of the exemplary computing environment are illustrated herein as discrete units (e.g., OS **51** being stored on non-volatile data storage device **51** and loaded into system memory **35** for use) such processes and components may reside or be processed at various times in different components of computing device **10**, remote computing devices **80**, and/or cloud-based services **90**.

(259) Remote computing devices **80** are any computing devices not part of computing device **10**. Remote computing devices **80** include, but are not limited to, personal computers, server computers, thin clients, thick clients, personal digital assistants (PDAs), mobile telephones, watches, tablet computers, laptop computers, multiprocessor systems, microprocessor based systems, set-top boxes, programmable consumer electronics, video game machines, game consoles, portable or handheld gaming units, network terminals, desktop personal computers (PCs), minicomputers, main frame computers, network nodes, and distributed or multi-processing computing environments. While remote computing devices **80** are shown for clarity as being separate from cloud-based services **90**, cloud-based services **90** are implemented on collections of networked remote computing devices **80**.

(260) Cloud-based services **90** are Internet-accessible services implemented on collections of networked remote computing devices **80**. Cloud-based services are typically accessed via application programming interfaces (APIs) which are software interfaces which provide access to computing services within the cloud-based service via API calls, which are pre-defined protocols for requesting a computing service and receiving the results of that computing service. While cloud-based services may comprise any type of computer processing or storage, three common categories of cloud-based services **90** are microservices **91**, cloud computing services **92**, and distributed computing services **93**.

(261) Microservices **91** are collections of small, loosely coupled, and independently deployable computing services. Each microservice represents a specific computing functionality and runs as a separate process or container. Microservices promote the decomposition of complex applications into smaller, manageable services that can be developed, deployed, and scaled independently. These services communicate with each other through well-defined application programming interfaces (APIs), typically using lightweight protocols like HTTP or message queues. Microservices **91** can be combined to perform more complex processing tasks.

(262) Cloud computing services **92** are delivery of computing resources and services over the Internet **75** from a remote location. Cloud computing services **92** provide additional computer hardware and storage on as-needed or subscription basis. Cloud computing services **92** can provide large amounts of scalable data storage, access to sophisticated software and powerful server-based processing, or entire computing infrastructures and platforms. For example, cloud computing services can provide virtualized computing resources such as virtual machines, storage, and networks, platforms for developing, running, and managing applications without the complexity of infrastructure management, and complete software applications over the Internet on a subscription basis.

(263) Distributed computing services **93** provide large-scale processing using multiple interconnected computers or nodes to solve computational problems or perform tasks collectively. In distributed computing, the processing and storage capabilities of multiple machines are leveraged to work together as a unified system. Distributed computing services are designed to

address problems that cannot be efficiently solved by a single computer or that require large-scale computational power. These services enable parallel processing, fault tolerance, and scalability by distributing tasks across multiple nodes.

(264) Although described above as a physical device, computing device **10** can be a virtual computing device, in which case the functionality of the physical components herein described, such as processors **20**, system memory **30**, network interfaces **40**, and other like components can be provided by computer-executable instructions. Such computer-executable instructions can execute on a single physical computing device, or can be distributed across multiple physical computing devices, including being distributed across multiple physical computing devices in a dynamic manner such that the specific, physical computing devices hosting such computer-executable instructions can dynamically change over time depending upon need and availability. In the situation where computing device **10** is a virtualized device, the underlying physical computing devices hosting such a virtualized computing device can, themselves, comprise physical components analogous to those described above, and operating in a like manner. Furthermore, virtual computing devices can be utilized in multiple layers with one virtual computing device executing within the construct of another virtual computing device. Thus, computing device **10** may be either a physical computing device or a virtualized computing device within which computer-executable instructions can be executed in a manner consistent with their execution by a physical computing device. Similarly, terms referring to physical components of the computing device, as utilized herein, mean either those physical components or virtualizations thereof performing the same or equivalent functions.

(265) The skilled person will be aware of a range of possible modifications of the various aspects described above. Accordingly, the present invention is defined by the claims and their equivalents.

Claims

1. A system for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks, comprising: a computing system comprising at least a memory and a processor; two or more datasets that are substantially correlated and which have been compressed with lossy compression, the two or more datasets comprising genomic data and proteomic data; a deep learning neural network configured to recover lost information associated with a compressed bit stream; a decoder comprising a first plurality of programming instructions that, when operating on the processor, cause the computing system to: receive a compressed bit stream, the compressed bit stream comprising cross-correlated genomic data and proteomic data; and decompress each of the compressed bit stream; and use the decompressed bit stream as an input into the deep learning neural network to recover lost information associated with the genomic data and proteomic data; a proteomic data upsampling neural network configured to upsample compressed proteomic data; a protein structure prediction model configured to predict 3D protein structures from upsampled proteomic data; a binding site prediction model configured to predict binding sites using the predicted 3D protein structures and upsampled proteomic data; a post-processing module configured to format, filter, validate, and error-check the output data; and a deep learning system with a latent transformer core for large codeword models, the deep learning system configured to: receive a plurality of input vectors; generate a plurality of latent space vectors by processing the plurality of input vectors through a variational autoencoder's encoder; learn relationships between the plurality of latent space vectors by processing the plurality of latent space vectors through a transformer, wherein the transformer does not include an embedding layer and a positional encoding layer; use the learned relationships between the plurality of latent space vectors to generate a plurality of output latent space vectors based on the plurality of input vectors; and generate output vectors by passing the plurality of output latent space vectors through the variational autoencoder's decoder.

2. The system of claim 1, wherein the genomic data comprises parallel genome datasets.
3. The system of claim 1, wherein the two or more datasets comprise genomic data from a subset of the human genome.
4. The system of claim 1, wherein the deep learning neural network is a neural network that can recover signals from a compressed bitstream.
5. The system of claim 1, wherein the compressed bit stream comprises a plurality of channels, wherein each of the plurality of channels is associated with a genomic dataset or a proteomic dataset.
6. The system of claim 1, further comprising a multi-modal protein data graph representation generator configured to create a multi-modal graph representation integrating proteomic data, predicted structures, and binding sites.
7. The system of claim 1, wherein the input vectors may contain a plurality of appended zeros and a plurality of truncated data points may be used to train and operationalize a transformer that predicts the next sequential vector following an input vector.
8. The system of claim 1, wherein the input vectors may contain a plurality of appended metadata.
9. The system of claim 8, wherein metadata comprises data type, temporal information, data source, data characteristics, and domain-specific metadata.
10. The system of claim 1, wherein the plurality of input vectors comprises a plurality of codewords.
11. The system of claim 10, wherein the plurality of codewords are converted into the plurality of latent space vectors.
12. A method for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks, comprising the steps of: training a deep learning neural network to recover lost information associated with a compressed bit stream; receiving the compressed bit stream, the compressed bit stream comprising cross-correlated genomic data and proteomic data; decompressing the compressed bit stream; using the decompressed bit stream as an input into the deep learning neural network to recover information lost during lossy compression of the genomic data and proteomic data; upsampling compressed proteomic data using a proteomic data upsampling neural network; predicting 3D protein structures from the upsampled proteomic data using a protein structure prediction model; predicting binding sites using the predicted 3D protein structures and upsampled proteomic data using a binding site prediction model; post-processing the output data through formatting, filtering, validation, and error handling steps; receiving a plurality of input vectors; generating a plurality of latent space vectors by processing the plurality of input vectors through a variational autoencoder's encoder; learning relationships between the plurality of latent space vectors by processing the plurality of latent space vectors through a transformer, wherein the transformer does not include an embedding layer and a positional encoding layer; using the learned relationships between the plurality of latent space vectors to generate a plurality of output latent space vectors based on the plurality of input vectors; and generating output vectors by passing the plurality of output latent space vectors through the variational autoencoder's decoder.
13. The method of claim 12, further comprising the step of generating a multi-modal protein data graph representation to integrate and visualize relationships between proteomic data, predicted structures, and binding sites.
14. The method of claim 12, wherein the genomic data comprises parallel genome datasets.
15. The method of claim 12, wherein the two or more datasets comprise genomic data from a subset of the human genome.
16. The method of claim 12, wherein the deep learning neural network is a neural network that can recover signals from a compressed bitstream.
17. The method of claim 12, wherein the compressed bit stream comprises a plurality of channels, wherein each of the plurality of channels is associated with a genomic dataset or a proteomic

dataset.

18. The method of claim 12, wherein the input vectors may contain a plurality of appended zeros and a plurality of truncated data points may be used to train and operationalize a transformer that predicts the next sequential vector following an input vector.

19. The method of claim 12, wherein the input vectors may contain a plurality of appended metadata.

20. The method of claim 19, wherein metadata comprises data type, temporal information, data source, data characteristics, and domain-specific metadata.

21. The method of claim 12, wherein the plurality of input vectors comprises a plurality of codewords.

22. The method of claim 21, wherein the plurality of codewords are converted into the plurality of latent space vectors.

23. One or more non-transitory computer-storage media having computer-executable instructions embodied thereon that, when executed by one or more processors of a computing system, cause the computing system to perform a method for upsampling of decompressed genomic and proteomic data after lossy compression using neural networks, the method comprising the steps of: training a deep learning neural network to recover lost information associated with a compressed bit stream; receiving the compressed bit stream, the compressed bit stream comprising cross-correlated genomic data and proteomic data; decompressing the compressed bit stream; using the decompressed bit stream as an input into the deep learning neural network to recover information lost during lossy compression of the genomic data and proteomic data; upsampling compressed proteomic data using a proteomic data upsampling neural network; predicting 3D protein structures from the upsampled proteomic data using a protein structure prediction model; predicting binding sites using the predicted 3D protein structures and upsampled proteomic data using a binding site prediction model; post-processing the output data through formatting, filtering, validation, and error handling steps; receiving a plurality of input vectors; generating a plurality of latent space vectors by processing the plurality of input vectors through a variational autoencoder's encoder; learning relationships between the plurality of latent space vectors by processing the plurality of latent space vectors through a transformer, wherein the transformer does not include an embedding layer and a positional encoding layer; using the learned relationships between the plurality of latent space vectors to generate a plurality of output latent space vectors based on the plurality of input vectors; and generating output vectors by passing the plurality of output latent space vectors through the variational autoencoder's decoder.

24. A system for integrated analysis of genomic and proteomic data, comprising: a data compression module configured to perform lossy compression on genomic and proteomic data; a neural upsampling module configured to restore compressed genomic and proteomic data; a structure prediction module configured to predict 3D protein structures from upsampled proteomic data; a binding site prediction module configured to identify potential binding sites using the predicted 3D structures and upsampled data; a data integration module configured to generate a graph-based multi-modal protein data graph representation of relationships between genomic data, proteomic data, predicted structures, and binding sites; a post-processing module configured to format, filter, validate, and error-check the output data.

25. The system of claim 24, wherein the multi-modal protein data graph representation comprises: protein nodes representing individual proteins; peptide nodes representing identified peptides; modification nodes representing post-translational modifications; structure nodes representing predicted 3D structures; binding site nodes representing predicted binding sites; edges connecting related nodes and capturing relationships between different data types.
