

US Patent & Trademark Office

Patent Public Search | Text View

United States Patent	12417388
Kind Code	B1
Date of Patent	September 16, 2025
Inventor(s)	Nataraj; Bindiganavale S.

Hardware implementations of activation functions in neural networks

Abstract

Circuitry for performing neural-network calculations includes a plurality of compute circuits, arranged in parallel with respective inputs and outputs, to receive function arguments for a node of a neural network on their respective inputs, compute values of a plurality of activation functions using the function arguments, and provide the values on their respective outputs. Each compute circuit of the plurality of compute circuits is to compute the values of a respective activation function of the plurality of activation functions. The circuitry also includes a multiplexor to select between the respective outputs of the plurality of compute circuits and to provide the values on a selected output as activation-function values for the node of the neural network, based on an activation-function selection signal.

Inventors:	Nataraj; Bindiganavale S. (Cupertino, CA)
Applicant:	Nataraj; Bindiganavale S. (Cupertino, CA)
Family ID:	1000005226470
Appl. No.:	17/083186
Filed:	October 28, 2020

Related U.S. Application Data

us-provisional-application US 63034907 20200604

Publication Classification

Int. Cl.: G06N3/084 (20230101); G06N3/04 (20230101)

U.S. Cl.:

CPC G06N3/084 (20130101); G06N3/04 (20130101);

Field of Classification Search

CPC: G06N (3/084); G06N (3/04); G06N (3/08); G06N (3/06); G06N (3/063)

References Cited

U.S. PATENT DOCUMENTS

Patent No.	Issued Date	Patentee Name	U.S. Cl.	CPC
2017/0103321	12/2016	Henry	N/A	G06N 3/04
2019/0325289	12/2018	Cammarota	N/A	G06N 3/04
2020/0074293	12/2019	Chin	N/A	G06N 3/04
2020/0327271	12/2019	Langhammer	N/A	G06N 3/063
2021/0232899	12/2020	Takamaeda	N/A	G06N 3/045
2021/0304010	12/2020	Sengupta	N/A	G06N 3/04

Primary Examiner: Afshar; Kamran

Assistant Examiner: Hales; Brian J

Attorney, Agent or Firm: Huse IP Law

Background/Summary

RELATED APPLICATION (1) This application claims priority to U.S. Provisional Patent Application No. 63/034,907, filed on Jun. 4, 2020, which is incorporated by reference in its entirety.

TECHNICAL FIELD

(1) This disclosure relates to neural networks, and more specifically to computing activation functions in neural networks.

BACKGROUND

(2) Neural networks use activation functions repeatedly, both for training and in deployment. Activation functions are traditionally computed by software running on a server in the cloud.

SUMMARY

(3) There is a need for techniques for computing activation functions in a quick, computationally efficient, and low-power manner.

(4) In some embodiments, circuitry includes a plurality of compute circuits, arranged in parallel and having respective inputs and outputs, to receive function arguments for a node of a neural network on their respective inputs, compute values of a plurality of activation functions using the function arguments, and provide the values on their respective outputs. Each compute circuit of the plurality of compute circuits is to compute the values of a respective activation function of the plurality of activation functions. The circuitry also includes a multiplexor to select between the respective outputs of the plurality of compute circuits and to provide the values on a selected output as activation-function values for the node of the neural network, based on an activation-function selection signal.

(5) In some embodiments, a method of performing neural-network calculations includes providing a function argument for a node of a neural network to one or more compute circuits of a plurality of compute circuits. The plurality of compute circuits is arranged in parallel and has respective inputs

and outputs. Each compute circuit of the plurality of compute circuits is configured to compute a value of a respective activation function of a plurality of activation functions using the function argument. The method also includes selecting an output of a respective compute circuit of the plurality of compute circuits, based on an activation-function selection signal, and providing a value on the selected output as an activation-function value for the node of the neural network.

Description

BRIEF DESCRIPTION OF THE DRAWINGS

- (1) For a better understanding of the various described implementations, reference should be made to the Detailed Description below, in conjunction with the following drawings.
- (2) FIG. 1 shows a neural network.
- (3) FIG. 2 is a functional block diagram of a neuron in a neural network.
- (4) FIG. 3 is a block diagram of circuitry to perform computations for a node in a neural network, in accordance with some embodiments.
- (5) FIG. 4 is a block diagram of an example of the circuitry of FIG. 3, in accordance with some embodiments.
- (6) FIG. 5 shows neural-network training for a deep neural network in accordance with some embodiments.
- (7) FIG. 6 shows an example of the neural-network training of FIG. 5, in accordance with some embodiments.
- (8) FIG. 7 is a flowchart showing a method of performing neural-network computations in accordance with some embodiments.
- (9) Like reference numerals refer to corresponding parts throughout the drawings and specification.

DETAILED DESCRIPTION

- (10) Reference will now be made in detail to various embodiments, examples of which are illustrated in the accompanying drawings. In the following detailed description, numerous specific details are set forth in order to provide a thorough understanding of the various described embodiments. However, it will be apparent to one of ordinary skill in the art that the various described embodiments may be practiced without these specific details. In other instances, well-known methods, procedures, components, circuits, and networks have not been described in detail so as not to unnecessarily obscure aspects of the embodiments.
- (11) FIG. 1 shows a neural network **100**. The neural network **100** includes a plurality of nodes **110** divided into layers. The nodes **110** may also be referred to as neurons. The layers include an input layer **102**, first hidden layer **104**, second hidden layer **106**, and output layer **108**. Each layer may have multiple nodes **110**. The neural network **100** as shown in FIG. 1 is an example of a deep neural network, because it has multiple hidden layers between the input layer **102** and output layer **108**. In some embodiments, the neural network **100** has more than two hidden layers (e.g., has three hidden layers). In some embodiments, the neural network **100** only has a single hidden layer, such that it is not a deep neural network. The nodes **110** of the input layer **102** receive input values **112** (i.e., input data) on respective inputs. The input values **112** are the input values for the neural network **100**. The nodes **110** of the following layers have inputs that receive input values (i.e., input data) from the nodes **110** of previous layers. For example, the nodes **110** of the first hidden layer **104** receive input values from the nodes **110** (e.g., each node **110**) of the input layer **102**, the nodes **110** of the second hidden layer **106** receive input values from the nodes **110** (e.g., each node **110**) of the first hidden layer **104**, and the nodes **110** of the output layer **108** receive input values from the nodes **110** (e.g., each node **110**) of the second hidden layer **106**. The nodes **110** of the output layer **108** provide respective output values **114** on respective outputs. The output values **114** are the output values for the neural network **100**.

(12) FIG. 2 is a functional block diagram of a neuron **200** in a neural network. The neuron **200** may be an example of each node **110** in the neural network **100** (FIG. 1). Assuming the neuron **200** is in a hidden layer (e.g., hidden layer **104** or **106**, FIG. 1) or output layer (e.g., output layer **108**, FIG. 1), the neuron **200** receives a vector X that includes values from the outputs of the neurons (e.g., each neuron) in a previous layer of the neural network to which the inputs of the neuron **200** are coupled. (If the neuron **200** is in the input layer **102**, it may have a single input **102** that receives input values.) A multiplication module **202** multiplies the vector X by a vector of weights W , to produce the product $W \cdot X$. Each weight in the vector W is for a respective input of the neuron **200**, and thus corresponds to (i.e., weights) values from a respective neuron in the previous layer of the neural network. An addition module **204** adds a bias b to the product $W \cdot X$, to produce the sum $W \cdot X + b$. The multiplication module **202** and the addition module **204** compose a module **206** that receives the vector X and provides $W \cdot X + b$. An activation-function module **208** computes an activation function $f(x)$, using $W \cdot X + b$, as received from the module **206**, as the argument x . The value of the activation function $f(x)$, as computed using $W \cdot X + b$ as the argument x , is provided on an output for the neuron **200**.

(13) The activation function $f(x)$ may be linear or non-linear. Different layers in a neural network may use different activation functions. Some layers in a neural network may use the same activation functions, while other layers in the neural network may use a different activation function.

(14) A neural network (e.g., the neural network **100**, FIG. 1) (e.g., with neurons **200** as nodes) may operate in two modes: a training mode and an operating mode. The operating mode is used after the neural network has been trained and deployed. During the training mode (e.g., during neural-network training **500**, FIG. 5, such as neural-network training **600**, FIG. 6), weights and biases are determined iteratively, by comparing outputs of the neural network to expected outputs and updating the weights and biases accordingly. During the operating mode, the weights and biases remain fixed, and the neural network provides outputs based on given inputs. A layer may use the same activation function in the operating mode as in the training mode, or may use a different activation function in the operating mode than in the training mode. The activation function used for a layer may be changed in the operating mode (i.e., once the neural network has been deployed). For example, after training, a neural network may be deployed such that a particular layer uses a particular activation function in the operating mode. The neural network may subsequently be changed so that the particular layer uses a different activation function in the operating mode. The weights and bias, however, remain unchanged; they are only changed during the training mode.

(15) FIG. 3 is a block diagram of circuitry **300** to perform computations for a node **110** (FIG. 1) (e.g., for the neuron **200**, FIG. 2) in accordance with some embodiments. The circuitry **300** includes a plurality of compute circuits **308-1** through **308-n** arranged in parallel. Each of the compute circuits **308** is configured to compute values of a respective activation function $f(x)$. For example, a first compute circuit **308-1** is configured to compute values of a first activation function $f_{\text{sub.1}}(x)$, a second compute circuit **308-2** is configured to compute values of a second activation function $f_{\text{sub.2}}(x)$, a third compute circuit **308-3** is configured to compute values of a third activation function $f_{\text{sub.3}}(x)$, and an n th compute circuit **308-n** is configured to compute values of an n th activation function $f_{\text{sub.n}}(x)$, where n is an integer greater than one. Each of the compute circuits **308** has a respective input **306** and a respective output **310**. The inputs **306** receive function arguments, which the compute circuits **308** use as arguments to compute respective activation-function values. The respective activation-function values are provided on respective outputs **310**.

(16) The activation functions computed by the compute circuits **308** may be linear and/or non-linear functions. Examples of activation functions that may be computed by respective compute circuits **308** include, without limitation, a sigmoid function, a hyperbolic tangent function, a rectified linear unit (ReLU) function, a leaky ReLU function, a max-pooling function, an average-

pooling function, and a zero-activation function. The max-pooling and average-pooling functions are used to down-sample matrices by selecting maximum or average values from portions of the matrices. The zero-activation function provides all zeros. A respective compute circuit **308** (e.g., as designed using register-transfer level (RTL) digital logic) may include a multiply-accumulate unit to generate an activation function (e.g., a non-linear activation function) by calculating a Taylor-series approximation of the activation function. For example, a first compute circuit **308** may include a multiply-accumulate unit to generate a sigmoid function by calculating its Taylor-series approximation, and/or a second compute circuit **308** may include a multiply-accumulate unit to generate a hyperbolic tangent function by calculating its Taylor-series approximation.

(17) FIG. **4** is a block diagram of circuitry **400** that is an example of the circuitry **300** (FIG. **3**) in accordance with some embodiments. The circuitry **400** includes a compute circuit **402** to compute a zero-activation function, a compute circuit **404** to compute a sigmoid function, a compute circuit **406** to compute a ReLU function, a compute circuit **408** to compute a hyperbolic tangent ($\tanh$) function, and a compute circuit **410** to compute a max-pooling ($\max$) function. The compute circuits **402**, **404**, **406**, **408**, and **410** are an example of the compute circuits **308-1** through **308-n** (FIG. **3**) (in this example, n equals 5) and are arranged in parallel. Numerous other examples of the compute circuits **308-1** through **308-n** are possible.

(18) The circuitry **300** (FIG. **3**) (e.g., the circuitry **400**, FIG. **4**) also includes a multiplexor **312**, which has inputs connected to the outputs **310**. The multiplexor **312** selects between the outputs **310** of the compute circuits **308** and provides the values on the selected output **310** as activation-function values for the node **110**, based on an activation-function selection signal **316**. The activation-function selection signal **316** may be provided to a function-selector circuit **318**, coupled to the multiplexor **312**, which forward the activation-function selection signal **316** or provides a corresponding signal to the multiplexor **312**. The activation-function selection signal **316** specifies the output **310** to be selected by the multiplexor **312**, and thereby specifies the activation function to be used for the node **110**.

(19) In some embodiments, the circuitry **300** (e.g., the circuitry **400**, FIG. **4**) further includes a pre-processor **304** that provides the function arguments to the respective inputs **306** of the plurality of compute circuits **308**. The compute circuits **308** are arranged in parallel between the pre-processor **304** and the multiplexor **312**. The pre-processor **304** may provide a particular function argument to each input **306** of the plurality of compute circuits **308**, which may operate simultaneously to produce respective values. The multiplexor **312**, by selecting a respective output **310**, effectively selects the compute circuit **308** that provides the result for the node **110** (i.e., provides for the node **110** the activation-function value corresponding to the particular function argument).

(20) Alternatively, the pre-processor **304** may include a de-multiplexor that provides a respective function argument to a respective input **306** of a selected compute circuit **308**, based on the activation-function selection signal **316**. The function-selector circuit **318** may be coupled to the de-multiplexor to forward the activation-function selection signal **316**, or a corresponding signal, to the de-multiplexor. The de-multiplexor does not provide the respective function argument to inputs **306** of the other, unselected compute circuits **308**, which may be disabled (e.g., placed in a low-power state) while the selected compute circuit **308** computes an activation-function value using the respective function argument. The multiplexor **312** selects the output **310** of the selected compute circuit **308** based on the activation-function selection signal **316** (or a corresponding signal) and provides the activation-function value computed by the selected compute circuit **308** as the activation-function value for the node **110**.

(21) The pre-processor **304** may receive input data **302** (i.e., input values) for the node **110** on an input **303** and may include processor circuitry **305** to generate the function arguments using the input data **302**, weights for the node **110**, and a bias for the node **110**. The input data **302** may come from the nodes **110** in a preceding layer of the neural network. For example, the processor circuitry **305** implements the functionality of the module **206** (FIG. **2**): it calculates $W \cdot \text{Math} \cdot X + b$, given

vector **X** as the input data **302**, and provides $W \cdot \text{Math.X} + b$ to one or more of the inputs **306** (e.g., to the input **306** of a selected compute circuit **308**, or to all of the inputs **306**). In some embodiments, the processor circuitry **305** includes a processor core (e.g., a central-processing-unit (CPU) core, graphics-processing-unit (GPU) core, or microcontroller core) and memory (e.g., non-volatile memory that serves as a non-transitory computer-readable medium) storing instructions (e.g., one or more programs) for execution by the processor core to calculate the function arguments. The memory may be embedded in and/or separate from the processor core. In some embodiments, the processor circuitry **305** includes an arithmetic logic unit (ALU) and associated state machine.

(22) In some embodiments, the circuitry **300** (e.g., the circuitry **400**, FIG. 4) includes a cache memory **314** coupled to the multiplexor **312** (e.g., to the output of the multiplexor **312**) and the pre-processor **304**. The cache memory **314** may store activation-function values. For example, activation-function values (e.g., as received from the multiplexor **312**) may be stored in the cache memory **314** in association with indicators of respective neurons **110** and/or layers for the activation-function values. The cache memory **314** may also, or alternatively, store the weights and bias for the node **110**. For example, the pre-processor **304** (e.g., the processor circuitry **305**) may cache the weights and the bias (e.g., W and b) for the node **110** in the cache memory **314** and retrieve the weights and the bias from the cache memory **314** when needed (e.g., to implement the functionality of module **206**, FIG. 2). The pre-processor **304** (e.g., the processor circuitry **305**) may update the weights and the bias during training (e.g., during backward propagation **514**, as described below for FIGS. 5 and 6) and cache the updated weights and the updated bias in the cache memory **314**. The circuitry **300** may be used for multiple nodes **110** (e.g., neurons **200**), and the weights and bias for each node **110** may be cached in the cache memory **314** and retrieved from the cache memory **314** to perform computations for that node **110**.

(23) The cache memory **314** may be used to implement skip connections in the neural network. In a skip connection, an activation-function value from a layer (e.g., layer m , where m is an integer indexing the layer) bypasses the next layer (i.e., layer $m+1$) and is provided to a subsequent layer (e.g., the layer following the next layer, which is layer $m+2$). This activation-function value may be cached in the cache **314** and then retrieved from the cache **314** as one of the input values provided to the subsequent layer.

(24) In some embodiments, the circuitry **300** is used for each node **110** in a layer (e.g., input layer **102**, hidden layer **104** or **106**, or output layer **108**, FIG. 1) of a neural network (i.e., for all of the nodes **110** in the layer). The layer has a plurality of nodes **110**, and the circuitry **300** is used for all of those nodes **110**. The pre-processor **304** provides function arguments for the plurality of nodes **110** in the layer to one or more compute circuits **308** (e.g., to a single compute circuit **308**, or to all of the compute circuits **308**) of the plurality of compute circuits **308-1** through **308-n**. The processor circuit **305** may generate the function arguments for all of the nodes **110** in the layer.

(25) In some embodiments, the circuitry **300** is used for a subset of the nodes **110** in a layer (e.g., input layer **102**, hidden layer **104** or **106**, or output layer **108**, FIG. 1) of a neural network. The layer has a plurality of nodes **110**, and the circuitry **300** is used for a subset (e.g., two or more) of those nodes **110**. The pre-processor **304** provides function arguments for the subset to one or more compute circuits **308** (e.g., to a single compute circuit **308**, or to all of the compute circuits **308**) of the plurality of compute circuits **308-1** through **308-n**. The processor circuit **305** may generate the function arguments for the subset.

(26) In some embodiments, the circuitry **300** is used for each node **110** in two or more layers (e.g., input layer **102**, hidden layer **104** and/or **106**, and/or output layer **108**, FIG. 1) (e.g., in every layer) of a neural network (i.e., for all of the nodes in each of the two or more layers) (e.g., for each node in every layer). Each of the two or more layers has a respective plurality of nodes **110**, and the circuitry **300** is used for all of those nodes **110**. The pre-processor **304** provides function arguments for the respective pluralities of nodes **110** in the two or more layers to one or more compute circuits **308** (e.g., to a single compute circuit **308** for each layer, or to all of the compute circuits **308**) of the

plurality of compute circuits **308-1** through **308-n**. The processor circuit **305** may generate the function arguments for all of the nodes **110** in the two or more layers.

(27) FIG. 5 shows neural-network training **500** for a deep neural network (e.g., neural network **100**, FIG. 1) in accordance with some embodiments. (Similar training may be performed for neural networks that are not deep.) The deep neural network includes an input layer **502** (e.g., input layer **102**, FIG. 1), hidden layers **504** and **506** (e.g., hidden layers **104** and **106**, FIG. 1), and output layer **508** (e.g., output layer **108**, FIG. 1). Each of these layers has a plurality of nodes **110** (FIG. 1) (e.g., neurons **200**, FIG. 2), which are not shown in FIG. 5 for simplicity. The neural-network training **500** includes two procedures, forward propagation **510** and backward propagation **514**, that are performed repeatedly (e.g., in an alternating manner). The neural-network training uses a training data set that includes expected output values of the neural network for respective input values **112**.

(28) During forward propagation **510**, the input layer **502** receives input values **112** and uses them to generate output values that are provided to the hidden layer **504**. The hidden layer **504**, hidden layer **506**, and output layer **508** repeat this process in turn, receiving the previous layer's output values as input values and generating respective output values (e.g., as described for the neural network **100**, FIG. 1). In generating respective output values, the input layer **502** calls an activation function **512-1**, the hidden layer **504** calls an activation function **512-2**, the hidden layer **506** calls an activation function **512-3**, and the output layer **508** calls an activation function **512-4**. To call each activation function, function arguments are provided to the circuitry **300** (FIG. 3), or to a particular instance of the circuitry **300**. In some embodiments, two or more of the activation functions **512-1** through **512-4** are the same function. For example, all of the layers use the same activation function except for the output layer **508**, which uses a distinct activation function **512-4** (i.e., the activation functions **512-1**, **512-2**, and **512-3** are the same function, which is distinct from the activation function **512-4**). The output values (i.e., activation-function values) generated by the output layer **508** during forward propagation **510** are cached in the cache memory **314** (FIGS. 3-4), in accordance with some embodiments. For example, the multiplexor **312** (FIGS. 3-4) provides the output values from the output layer **508** to the cache memory **314**.

(29) The output values from the output layer **508** for forward propagation **510** are compared (**516**) to the expected output values, and backward propagation **514** is performed using the results of this comparison. A difference value, which corresponds to the difference between the expected and actual output values, is provided to the output layer **508**. The output layer **508** retrieves its activation-function values from forward propagation, along with its weights and biases, from the cache memory **314** (FIG. 3). A loss function is applied to the activation-function values in accordance with the difference value, and the weights and biases are updated accordingly. The updated weights and biases are cached in the cache memory **314** for use during the next round of forward propagation **510**. This process is repeated for the hidden layers **506** and **504** until the weights and biases of the output layer **508** and hidden layers **506** and **504** have been updated and cached in the cache memory **314**, at which point this round of backward propagation **514** is complete. The next round of forward propagation **510** (e.g., using a next group of input values **112** from the training set) is then performed, and the neural-network training **500** continues accordingly until the output values from the output layer **508** for forward propagation **510** converge with the expected output values.

(30) FIG. 6 shows neural-network training **600** that is an example of the neural-network training **500** (FIG. 5) in accordance with some embodiments. In the neural-network training **600**, the activation functions **512-1**, **512-2**, and **512-3** for the input layer **502**, hidden layer **504**, and hidden layer **506** are a ReLU function **602**. The activation function **512-4** for the output layer **508** is a sigmoid function **604**. The activation functions for the neural-network training **600** are merely one example of the activation functions **512-1** through **512-4**. Numerous other examples are possible.

(31) FIG. 7 is a flowchart showing a method **700** of performing neural-network computations in accordance with some embodiments. In the method **700**, a function argument for a node (e.g., a

node **110**, FIG. 1; a neuron **200**, FIG. 2) of a neural network (e.g., neural network **100**, FIG. 1) is provided (**702**) to one or more compute circuits of a plurality of compute circuits (e.g., compute circuits **308-1** through **308-n**, FIG. 3) (e.g., compute circuits **402**, **404**, **406**, **408**, and **410**, FIG. 4). The plurality of compute circuits is arranged in parallel and has respective inputs and outputs (e.g., inputs **306** and outputs **310**, FIGS. 3-4). Each compute circuit of the plurality of compute circuits is configured to compute a value of a respective activation function of a plurality of activation functions (e.g., activation functions **512-1** through **512-4**, FIG. 5) (e.g., activation functions **602** and **604**, FIG. 6) using the function argument. In some embodiments, a pre-processor (e.g., pre-processor **304**, FIGS. 3-4) coupled to the inputs of the plurality of compute circuits provides the function argument to respective inputs of the one or more compute circuits.

(32) In some embodiments, the function argument is generated (**704**) using input data (i.e., input values) for the node, weights for the node, and a bias for the node. For example, the pre-processor generates the function argument (e.g., in the module **206**, FIG. 2). The weights and the bias may be cached (**706**) in a cache memory (e.g., cache memory **314**, FIGS. 3-4).

(33) In some embodiments, the function argument is provided (**708**) to each input of the plurality of compute circuits in parallel, at the same time. Alternatively, a respective compute circuit of the plurality of compute circuits is selected (**710**) based on an activation-function selection signal (e.g., activation-function selection signal **316**, FIGS. 3-4), and the function argument is provided to the input of the selected compute circuit.

(34) An output of a respective compute circuit of the plurality of compute circuits is selected (**712**), based on the activation-function selection signal (e.g., activation-function selection signal **316**, FIGS. 3-4).

(35) A value on the selected output is provided (**714**) as an activation-function value for the node of the neural network. In some embodiments, the activation-function value for the node is cached (**716**) in the cache memory. The cache memory may be used to implement a skip connection for the activation-function value.

(36) In some embodiments, the activation-function value is retrieved (**718**) from the cache memory during backward propagation (e.g., backward propagation **514**, FIGS. 5-6) through the neural network in a training procedure for the neural network. The weights and the bias are updated (**720**) during the backward propagation. The updated weights and the updated bias are cached (**722**) in the cache memory.

(37) The method **700** (e.g., including providing the function argument in step **702**, selecting an output in step **712**, and providing the value on the selected output in step **714**) may be performed for some or all of the neurons in a neural network. For example, a respective instance of the method **700** may be performed for a respective node in a first layer of the neural network. The first layer includes a plurality of nodes, and additional instances of the method **700** may be performed for each additional node, or for some of the additional nodes, in the first layer. Respective instances of the method **700** (e.g., each including providing the function argument in step **702**, selecting an output in step **712**, and providing the value on the selected output in step **714**) thus may be performed for each node in the first layer or for a subset of the plurality of nodes in the first layer. In another example, respective instances of the method **700** (e.g., each including providing the function argument in step **702**, selecting an output in step **712**, and providing the value on the selected output in step **714**) are performed for each node in multiple layers of the neural network (e.g., for each node in the first layer and each node in a second layer) (e.g., with each of the multiple layers having multiple nodes).

(38) The circuitry **300** (FIG. 3) and the method **700** (FIG. 7) allow activation functions to be computed in a quick, computationally efficient, and low-power manner, by using dedicated compute circuits for specific activation functions. The circuitry **300** and the method **700** also can be used to compute multiple activation functions, since compute circuits for multiple activation functions are available in parallel. For example, the circuitry **300** and the method **700** may be used

to change an activation function used in a deployed neural network.

(39) The foregoing description, for purpose of explanation, has been described with reference to specific embodiments. However, the illustrative discussions above are not intended to be exhaustive or to limit the scope of the claims to the precise forms disclosed. Many modifications and variations are possible in view of the above teachings. The embodiments were chosen in order to best explain the principles underlying the claims and their practical applications, to thereby enable others skilled in the art to best use the embodiments with various modifications as are suited to the particular uses contemplated.

Claims

1. Circuitry for performing neural-network calculations, comprising: a plurality of compute circuits, arranged in parallel and having respective inputs and outputs, to receive function arguments for a node of a neural network on their respective inputs, compute values of a plurality of activation functions using the function arguments, and provide the values on their respective outputs, wherein each compute circuit of the plurality of compute circuits is to compute the values of a respective activation function of the plurality of activation functions; a pre-processor to provide the function arguments to the respective inputs of the plurality of compute circuits, the pre-processor comprising an input to receive input data for the node and processor circuitry to generate the function arguments using the input data, weights for the node, and a bias for the node; and a multiplexor to select between the respective outputs of the plurality of compute circuits and to provide the values on a selected output as activation-function values for the node of the neural network, based on an activation-function selection signal.
2. The circuitry of claim 1, wherein the pre-processor comprises a de-multiplexor to provide a respective function argument to a respective input of a selected compute circuit of the plurality of compute circuits, based on the activation-function selection signal.
3. The circuitry of claim 1, further comprising a cache memory, coupled to the multiplexor and the pre-processor, to store the activation-function values; wherein the pre-processor is to retrieve stored activation-function values from the cache memory during backward propagation through the neural network in a training procedure for the neural network.
4. The circuitry of claim 3, wherein the processor circuitry is to: cache the weights and the bias in the cache memory; retrieve the weights and the bias from the cache memory; update the weights and the bias during the backward propagation; and cache the updated weights and the updated bias in the cache memory.
5. The circuitry of claim 1, wherein: the node is a first node in a first layer of the neural network, the first layer comprising a plurality of nodes including the first node; and the pre-processor is to provide function arguments for the plurality of nodes of the first layer to one or more compute circuits of the plurality of compute circuits.
6. The circuitry of claim 1, wherein: the node is a first node in a first layer of the neural network, the first layer comprising a plurality of nodes including the first node; and the pre-processor is to provide function arguments for a subset of the plurality of nodes of the first layer to one or more compute circuits of the plurality of compute circuits, the subset comprising the first node.
7. The circuitry of claim 1, wherein: the node is a first node in a first layer of nodes of the neural network; the neural network comprises a plurality of layers of nodes, including the first layer and a second layer of nodes; and the pre-processor is to provide function arguments for the nodes of the first layer and the second layer to one or more compute circuits of the plurality of compute circuits.
8. The circuitry of claim 1, wherein the plurality of activation functions comprises activation functions selected from the group consisting of a sigmoid function, a hyperbolic tangent function, a rectified linear unit (ReLU) function, a leaky ReLU function, a max-pooling function, an average-pooling function, and a zero-activation function that provides all zeros.

9. The circuitry of claim 1, wherein: the plurality of activation functions comprises at least one of a sigmoid function or a hyperbolic tangent function; and the plurality of compute circuits comprises a compute circuit comprising a multiply-accumulate unit to calculate a Taylor-series approximation of the sigmoid function or the hyperbolic tangent function.
10. A method of performing neural-network calculations, comprising: using processor circuitry of a pre-processor, generating a function argument for a node of a neural network using input data for the node, weights for the node, and a bias for the node; providing the function argument for the node of the neural network from the pre-processor to one or more compute circuits of a plurality of compute circuits, wherein: the plurality of compute circuits is arranged in parallel and has respective inputs and outputs, the pre-processor is coupled to the inputs of the plurality of compute circuits, the pre-processor provides the function argument to respective inputs of the one or more compute circuits of the plurality of compute circuits, and each compute circuit of the plurality of compute circuits is configured to compute a value of a respective activation function of a plurality of activation functions using the function argument; selecting an output of a respective compute circuit of the plurality of compute circuits, based on an activation-function selection signal; and providing a value on the selected output as an activation-function value for the node of the neural network.
11. The method of claim 10, wherein providing the function argument from the pre-processor to the one or more compute circuits comprises providing the function argument from the pre-processor to each input of the plurality of compute circuits in parallel.
12. The method of claim 10, further comprising selecting a respective compute circuit of the plurality of compute circuits, based on the activation-function selection signal; wherein the pre-processor provides the function argument to the input of the selected compute circuit.
13. The method of claim 10, further comprising caching the activation-function value for the node in a cache memory.
14. The method of claim 13, further comprising retrieving the activation-function value from the cache memory during backward propagation through the neural network in a training procedure for the neural network.
15. The method of claim 14, further comprising: caching the weights and the bias in the cache memory; in the pre-processor, updating the weights and the bias during the backward propagation; and caching the updated weights and the updated bias in the cache memory.
16. The method of claim 10, wherein: the node is a respective node in a first layer of the neural network, the first layer comprising a plurality of nodes; and the method further comprises performing the providing of the function argument, the selecting of the output, and the providing of the value for each node in the first layer.
17. The method of claim 10, wherein: the node is a respective node in a first layer of the neural network, the first layer comprising a plurality of nodes; and the method further comprises performing the providing of the function argument, the selecting of the output, and the providing of the value for a subset of the plurality of nodes in the first layer.
18. The method of claim 10, wherein: the node is a respective node in a first layer of nodes of the neural network; the neural network comprises a plurality of layers of nodes, including the first layer and a second layer of nodes; and the method further comprises performing the providing of the function argument, the selecting of the output, and the providing of the value for each node in the first layer and each node in the second layer.
-