



US012417430B1

(12) **United States Patent**
Manyam et al.

(10) **Patent No.:** **US 12,417,430 B1**
(45) **Date of Patent:** **Sep. 16, 2025**

(54) **VERBALLY INTERACTIVE MATERIALS
HANDLING FACILITY**

(71) Applicant: **AMAZON TECHNOLOGIES, INC.,**
Reno, NV (US)

(72) Inventors: **Ohil Krishnamurthy Manyam,**
Bellevue, WA (US); **Rohit Prasad,**
Acton, MA (US)

(73) Assignee: **AMAZON TECHNOLOGIES, INC.,**
Reno, NV (US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 94 days.

(21) Appl. No.: **14/307,053**

(22) Filed: **Jun. 17, 2014**

(51) **Int. Cl.**

G06Q 30/00 (2023.01)
G06Q 10/087 (2023.01)
G10L 25/48 (2013.01)
G10L 15/065 (2013.01)

(52) **U.S. Cl.**

CPC **G06Q 10/087** (2013.01); **G10L 25/48**
(2013.01); **G06Q 30/00** (2013.01); **G10L**
15/065 (2013.01)

(58) **Field of Classification Search**

CPC G10L 15/065
USPC 704/270
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

7,225,980 B2 6/2007 Ku et al.
7,949,568 B2 5/2011 Fano et al.
8,009,864 B2 8/2011 Linaker et al.

8,189,855 B2 5/2012 Opalach et al.
8,630,924 B2 1/2014 Groenevelt et al.
9,235,928 B2 1/2016 Medioni et al.
9,785,952 B2 * 10/2017 Zinn G06Q 30/0256
2002/0107758 A1 * 8/2002 Takakura G06Q 30/02
705/26.7
2003/0113002 A1 * 6/2003 Philomin G06F 18/256
382/116
2004/0181461 A1 * 9/2004 Raiyani G06Q 30/0641
705/26.61
2005/0187758 A1 * 8/2005 Khasin G10L 15/005
704/10
2008/0294438 A1 * 11/2008 Geisel G01G 19/4144
704/251
2011/0011936 A1 1/2011 Morandi et al.
2012/0284132 A1 11/2012 Kim et al.
2013/0031476 A1 * 1/2013 Coin G06F 17/2881
715/706
2013/0284806 A1 10/2013 Margalit
2014/0201025 A1 * 7/2014 Adoni G06Q 30/0631
705/26.7

(Continued)

OTHER PUBLICATIONS

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994.*

(Continued)

Primary Examiner — Forrest F Tzeng


(74) *Attorney, Agent, or Firm* — Lindauer Law, PLLC

(57)

ABSTRACT

Described are systems and techniques configured to acquire data from user speech generated by one user of a materials handling facility speaking to themselves or another user. The user speech is processed using natural language processing to determine key words, concepts, meanings, and so forth. Based on the processed data, queries may be executed, data may be presented in user interfaces, operation of the materials handling facility may be modified, and so forth.

21 Claims, 11 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

2014/0257803	A1 *	9/2014	Yu	G10L 15/07 704/232
2014/0336537	A1 *	11/2014	Patel	A61B 5/0816 600/586
2014/0365334	A1 *	12/2014	Hurewitz	G06Q 30/0613 705/26.41
2015/0086107	A1	3/2015	Dedeoglu et al.	
2015/0235637	A1 *	8/2015	Casado	G10L 15/222 704/235
2015/0269642	A1 *	9/2015	Cai	G06Q 30/0613 705/14.49
2016/0259339	A1 *	9/2016	High	A47L 11/4011
2016/0292966	A1 *	10/2016	Denham	G06Q 30/0633

OTHER PUBLICATIONS

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994 (see previous Office action). (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994—see previous Office action attachment. (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994 - see previous Office action

attachment. - see attached reference in the previous Office action. (Year: 1994) (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994—(see attached reference in the previous Office action). (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994. (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance," Proceedings, IEEE Workshop on Mobile Computing Systems and Applications, 1994 (see attached reference in the previous Office actions). (Year: 1994).*

Asthana, et al., "An indoor wireless system for personalized shopping assistance", CiteSeerX, In Proceedings of IEEE Workshop on Mobile Computing Systems and Applications, 1994; [retrieved on Jun. 30, 2013]. Retrieved from the Internet: <URL:http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.127.3033>.

Kalnikaite, et al., "How to Nudge in Situ: Designing Lambent Devices to Deliver Information Salience in Supermarkets", ACM, In proceeding of: UbiComp 2011: Ubiquitous Computing, 13th International Conference, UbiComp 2011, Beijing, China, Sep. 17-21, 2011, Retrieved from Internet: <URL:http://www.researchgate.net/publication/221568350_How_to_nudge_in_Situ_designing_lambent_devices_to_deliver_salient_information_in_supermarkets>.

Pop, Christian, "Introduction to the BodyCom Technology", AN1391, DS01391A, Microchip Technology, Inc., May 2, 2011.

* cited by examiner




FIG. 1




FIG. 2




FIG. 3




FIG. 4




FIG. 5




FIG. 6




FIG. 7




FIG. 8




FIG. 9




FIG. 10

1100




FIG. 11

VERBALLY INTERACTIVE MATERIALS HANDLING FACILITY

BACKGROUND

Retailers, wholesalers, and other product distributors typically maintain an inventory of various items that may be ordered, purchased, leased, borrowed, rented, viewed, etc. by clients or customers. For example, an e-commerce website may maintain inventory in a fulfillment center. When a customer orders an item, the item is picked from inventory, routed to a packing station, packed and shipped to the customer. Likewise, physical stores maintain inventory in customer accessible areas (e.g., shopping area) and customers can pick items from inventory and take them to a cashier for purchase, rental, and so forth. Many of those physical stores also maintain inventory in a storage area, fulfillment center, or other facility that can be used to replenish inventory located in the shopping areas or to satisfy orders for items that are placed through other channels (e.g., e-commerce). Other examples of entities that maintain facilities holding inventory include libraries, museums, rental centers and the like. In each instance, for an item to be moved from one location to another, it is picked from its current location and transitioned to a new location. It is often desirable to provide to users information associated with the items in inventory, other users, or other information about operation of the facility.

BRIEF DESCRIPTION OF FIGURES

The detailed description is set forth with reference to the accompanying figures. In the figures, the left-most digit(s) of a reference number identifies the figure in which the reference number first appears. The use of the same reference numbers in different figures indicates similar or identical items or features.

FIG. 1 is a block diagram illustrating a materials handling facility configured to respond to verbal communication between a user of the facility and another person, according to some implementations.

FIG. 2 is a block diagram illustrating additional details of the materials handling facility, according to some implementations.

FIG. 3 illustrates a block diagram of a server configured to support operation of the facility, according to some implementations.

FIG. 4 is a block diagram of a tote, according to some implementations.

FIG. 5 illustrates a user device configured to acquire information about user context, present information, and so forth, according to some implementations.

FIG. 6 illustrates an overhead imaging sensor configured to acquire sensor data in the facility, according to some implementations.

FIG. 7 illustrates a user interface for setting how the facility interacts with the user, according to some implementations.

FIG. 8 illustrates user speech and the system responding to this speech, according to some implementations.

FIG. 9 illustrates the system responding to speech between two users of the facility, according to some implementations.

FIG. 10 depicts a flow diagram of a process for performing actions based on speech between users of the facility, according to some implementations.

FIG. 11 depicts a flow diagram of another process for performing actions based on speech between users of the facility, according to some implementations.

While implementations are described herein by way of example, those skilled in the art will recognize that the implementations are not limited to the examples or figures described. It should be understood that the figures and detailed description thereto are not intended to limit implementations to the particular form disclosed but, on the contrary, the intention is to cover all modifications, equivalents and alternatives falling within the spirit and scope as defined by the appended claims. The headings used herein are for organizational purposes only and are not meant to be used to limit the scope of the description or the claims. As used throughout this application, the word “may” is used in a permissive sense (i.e., meaning having the potential to), rather than the mandatory sense (i.e., meaning must). Similarly, the words “include,” “including,” and “includes” mean including, but not limited to.

DETAILED DESCRIPTION

This disclosure describes systems and techniques for providing information to a user of a materials handling facility (facility). The facility may include, or have access to, an inventory management system. The inventory management system may be configured to maintain information about items, users, condition of the facility, and so forth. For example, the inventory management system may maintain data indicative of what items a particular user is ordered to pick, historical data about items the particular user has picked, suggested routes within the facility to follow while picking, and so forth.

The systems and techniques described in this disclosure may be used to provide a user of the facility with information, modify operation of the facility or mechanisms therein, or otherwise assist the user. A plurality of imaging sensors configured to acquire image data, such as cameras, may also be arranged within the facility. The image data may be processed by the inventory management system to determine information such as identity of a user, location of the user, and so forth.

A plurality of microphones may be arranged within the facility and configured to acquire audio data representative of sound within the facility. The microphones may be fixed within the facility, mounted on totes, carried by the user, and so forth. In some implementations, a plurality of microphones may be configured in an array. Data acquired by the array may be processed using techniques such as beamforming to allow for localization of a sound. Localization of a sound includes determining a position in space from which the sound emanates. The data acquired by the microphones may be processed by a speech interaction system of the inventory management system.

The speech interaction system may be configured to process the audio data and determine a communicative sound issued from the user. The communicative sounds may include, but are not limited to, speech, non-lexical utterances such as laughs, sighs, snorts, fingersnaps, handclaps, and so forth. For ease of discussion and not necessarily by way of limitation, speech as used in this disclosure may include the communicative sounds described above. The speech interaction system may be configured to recognize speech to generate text or commands, voice recognition to identify a particular user based on the communicative sounds, and so forth. In some implementations the speech interaction system may use the image data or data from other sensors to

localize a source of sound. For example, the speech interaction system may compare facial movements with sounds to determine a particular user is speaking.

The speech interaction system is configured to process speech which is not directed to the speech interaction system. For example, the speech interaction system may process speech while a user is talking to themselves, speech when the user is talking to another person, speech when the user is talking to another system, and so forth. The other person or system to whom the user speaking may be physically present in the facility, may be in communication with the user by way of a cell phone or other audible communication mechanism, and so forth. For example, the user may be speaking to a personal information assistant application which executes at least in part on servers outside of the facility. The speech which is processed by the speech interaction system may not necessarily have a predefined syntax or grammar, instead comprising natural language. The speech interaction system may analyze attributes such as intonation, speaking rate, perplexity, and so forth to determine when speech is or is not directed to the speech interaction system. For example, based on an intonation which indicates a particular pattern of pitch, a speaking rate above 150 words per minute, and perplexity above a threshold value, the speech may be deemed to be directed at another user.

The speech interaction system determines one or more actions to perform based on the speech. These actions may include one or more of: retrieving information, storing information, generating a user interface for presentation to the user by way of an output device such as a speaker, moving inventory locations within the facility, moving inventory from one inventory location to another, requesting assistance from another user, and so forth. For example, the speech interaction system may hear the user say to another user "I can't find it" and determine that the user is unable to find a particular item to pick. The speech interaction system may then perform the actions of retrieving information indicating an inventory location of the particular item, generating user interface data including synthesized speech which is then presented using a speaker for the user to hear as an audible user interface. Continuing the example, synthesized speech may announce "The widget on your pick list is located ahead of you on rack A shelf 3." Thus, the user is presented with information which may be useful in performing their tasks in the facility, without having to initiate a query with the inventory management system.

The determination of the one or more actions to perform by the speech interaction system may be based at least in part on a user context. The user context provides information which may provide cues to the speech interaction system to help disambiguate words, determine meaning, select particular actions, and so forth. For example, the user context may include a direction of gaze of the user, proximity to other users, location of the user within the facility, and so forth. The user context may also include information which is specific to the user. For example, the user context may include historical activity such as a list of items picked since entry to the facility today.

The inventory management system may be configured to identify the user. Identification may comprise associating a user with a particular user account. For example, the identification may involve one or more of facial recognition, entry of one or more login credentials such as the username and password, detection of a tag receiving data stored within the tag such as a radiofrequency identification (RFID) chip, and so forth. Once identified, the user context may also

include user data which is specific to a particular identity. For example, the user data may include demographic data such as the user's age, relationship data indicative of a relationship between the identified user and another person, user preferences indicative of a preferred language, preferred wordbase, and so forth. Continuing the example, when the user's preferred language is English and the preferred wordbase is British English, audio data of the phrase "where is the torch" as added by the user may be correctly interpreted as an inquiry to find a flashlight item, and information indicating the location of the flashlight item may be presented in a user interface. In another example, the relationship data may provide additional meaning or context within which an appropriate action may be selected. For example, relationship data indicating that the user is a supervisor of a second user may result in interpretation of the phrase "could you please pick the widget from our overstock area" as an order to change the pick location for the item to the inventory location in the overstock area. In contrast, where the relationship data indicates that the user is a subordinate of the second user, the action may comprise presenting a user interface to the second user which includes a prompt to confirm the change in pick location.

The facility may include a materials handling facility, library, museum, and so forth. As used herein, a materials handling facility may include, but is not limited to, warehouses, distribution centers, cross-docking facilities, order fulfillment facilities, packaging facilities, shipping facilities, rental facilities, libraries, retail stores, wholesale stores, museums, or other facilities or combinations of facilities for performing one or more functions of materials (inventory) handling.

The systems and techniques described herein enable actions to be performed in the facility such as information presentation or movement of inventory in a way which is meaningful, unobtrusive, and responsive to the user without the user specifically requesting those actions. The speech interaction system may improve overall operation of the inventory management system, improve the experience of the users, and so forth.

Illustrative System

An implementation of a materials handling system **100** configured to store and manage inventory items is illustrated in FIG. 1. A materials handling facility **102** (facility) comprises one or more physical structures or areas within which one or more items **104(1)**, **104(2)**, . . . , **104(Q)** may be held. As used in this disclosure, letters in parenthesis such as "(Q)" indicate an integer value. The items **104** comprise physical goods, such as books, pharmaceuticals, repair parts, electronic gear, and so forth.

The facility **102** may include one or more areas designated for different functions with regard to inventory handling. In this illustration, the facility **102** includes a receiving area **106**, a storage area **108**, and a transition area **110**.

The receiving area **106** may be configured to accept items **104**, such as from suppliers, for intake into the facility **102**. For example, the receiving area **106** may include a loading dock at which trucks or other freight conveyances unload the items **104**.

The storage area **108** is configured to store the items **104**. The storage area **108** may be arranged in various physical configurations. In one implementation, the storage area **108** may include one or more aisles **112**. The aisle **112** may be configured with, or defined by, inventory locations **114** on one or both sides of the aisle **112**. The inventory locations **114** may include one or more of shelves, racks, cases, cabinets, bins, floor locations, or other suitable storage

5

mechanisms. The inventory locations **114** may be affixed to the floor or another portion of the facility's structure, or may be movable such that the arrangements of aisles **112** may be reconfigurable. In some implementations, the inventory locations **114** may be configured to move independently of an outside operator. For example, the inventory locations **114** may comprise a rack with a power source and a motor, operable by a computing device to allow the rack to move from one position within the facility **102** to another. Continuing the example, the inventory location **114** may move from one aisle **112** to another, from one position within an aisle **112** to another, and so forth.

One or more users **116** and totes **118** or other material handling apparatus may move within the facility **102**. For example, the user **116** may move about within the facility **102** to pick or place the items **104** in various inventory locations **114**, placing them on the tote **118** for ease of transport.

Instead of or in addition to the users **116** other mechanisms such as robots, forklifts, cranes, aerial drones, conveyors, elevators, pipes, and so forth, may move items **104** about the facility **102**. For example, a robot may pick the item **104** from a first inventory location **114(1)** and move it to a second inventory location **114(2)**.

The one or more users **116** may play different roles in the operation of the facility **102**. For example, some users **116** may be tasked with picking or placing items **104** while other users **116** support the operation of the facility **102** by assisting other users **116**. Furthermore, the roles of a particular user **116** may change over time. For example the user **116(1)** may start out as an assistant and later be tasked as a supervisor.

Different relationships may exist between users **116**. For example the user **116(1)** may be a supervisor of the user **116(2)**. In another example, the user **116(1)** may be a sibling, spouse, friend, or otherwise be related to the user **116(2)**.

One or more sensors **120** may be configured to acquire information in the facility **102**. The sensors **120** may include, but are not limited to, microphones, imaging sensors, weight sensors, radio frequency (RF) receivers, temperature sensors, humidity sensors, vibration sensors, and so forth. The sensors **120** may be stationary or mobile, relative to the facility **102**. For example, the inventory locations **114**, the totes **118**, or other devices such as user devices may contain microphones configured to acquire audio data. In another example, the inventory locations **114** may contain imaging sensors configured to acquire images of pick or placement of items **104** on shelves. The sensors **120** are discussed in more detail below with regard to FIG. 2.

During operation of the facility **102**, the sensors **120** may be configured to provide information suitable for tracking how objects move within the facility **102**. For example, a series of images acquired by an imaging sensor may indicate removal of an item **104** from a particular inventory location **114** by the user **116** and placement of the item **104** on or at least partially within the tote **118**. The tote **118** is discussed in more detail below with regard to FIG. 4.

While the storage area **108** is depicted as having one or more aisles **112**, inventory locations **114** storing the items **104**, sensors **120**, and so forth, it is understood that the receiving area **106**, the transition area **110**, or other areas of the facility **102** may be similarly equipped. Furthermore, the arrangement of the various areas within the facility **102** are depicted functionally rather than schematically. For example, in some implementations multiple different receiving areas **106**, storage areas **108**, and transition areas **110** may be interspersed rather than segregated.

6

The facility **102** may include, or be coupled to, an inventory management system **122**. The inventory management system **122** is configured to interact with users **116** or devices such as sensors **120**, robots, material handling equipment, computing devices, and so forth, in one or more of the receiving area **106**, the storage area **108**, or the transition area **110**.

The facility **102** may be configured to receive different kinds of items **104** from various suppliers, and to store them until a customer orders or retrieves one or more of the items **104**. A general flow of items **104** through the facility **102** is indicated by the arrows of FIG. 1. Specifically, as illustrated in this example, items **104** may be received from one or more suppliers, such as manufacturers, distributors, wholesalers, and so forth, at the receiving area **106**. In various implementations, the items **104** may include merchandise, commodities, perishables, or any suitable type of item, depending on the nature of the enterprise that operates the facility **102**.

Upon being received from a supplier at receiving area **106**, the items **104** may be prepared for storage. For example, in some implementations, items **104** may be unpacked or otherwise rearranged. The inventory management system **122** may include one or more software applications executing on a computer system to provide inventory management functions. These inventory management functions may include maintaining information indicative of the type, quantity, condition, cost, location, weight, or any other suitable parameters with respect to the items **104**. The items **104** may be stocked, managed, or dispensed in terms of countable, individual units or multiples, such as packages, cartons, crates, pallets, or other suitable aggregations. Alternatively, some items **104**, such as bulk products, commodities, and so forth, may be stored in continuous or arbitrarily divisible amounts that may not be inherently organized into countable units. Such items **104** may be managed in terms of measurable quantity such as units of length, area, volume, weight, time, duration, or other dimensional properties characterized by units of measurement. Generally speaking, a quantity of an item **104** may refer to either a countable number of individual or aggregate units of an item **104** or a measurable amount of an item **104**, as appropriate.

After arriving through the receiving area **106**, items **104** may be stored within the storage area **108**. In some implementations, like items **104** may be stored or displayed together in the inventory locations **114** such as in bins, on shelves, hanging from pegboards, and so forth. In this implementation, all items **104** of a given kind are stored in one inventory location **114**. In other implementations, like items **104** may be stored in different inventory locations **114**. For example, to optimize retrieval of certain items **104** having frequent turnover within a large physical facility, those items **104** may be stored in several different inventory locations **114** to reduce congestion that might occur at a single inventory location **114**.

When a customer order specifying one or more items **104** is received, or as a user **116** progresses through the facility **102**, the corresponding items **104** may be selected or "picked" from the inventory locations **114** containing those items **104**. In various implementations, item picking may range from manual to completely automated picking. For example, in one implementation, a user **116** may have a list of items **104** they desire and may progress through the facility **102** picking items **104** from inventory locations **114** within the storage area **108**, and placing those items **104** into a tote **118**. In other implementations, employees of the facility **102** may pick items **104** using written or electronic

pick lists derived from customer orders. These picked items **104** may be placed into the tote **118** as the employee progresses through the facility **102**.

After items **104** have been picked, they may be processed at a transition area **110**. The transition area **110** may be any designated area within the facility **102** where items **104** are transitioned from one location to another, or from one entity to another. For example, the transition area **110** may be a packing station within the facility **102**. When the item **104** arrives at the transition area **110**, the item **104** may be transitioned from the storage area **108** to the packing station. Information about the transition may be maintained by the inventory management system **122**.

In another example, if the items **104** are departing the facility **102** a list of the items **104** may be obtained and used by the inventory management system **122** to transition responsibility for, or custody of, the items **104** from the facility **102** to another entity. For example, a carrier may accept the items **104** for transport with that carrier accepting responsibility for the items **104** indicated in the list. In another example, a customer may purchase or rent the items **104** and remove the items **104** from the facility **102**.

The user **116** may produce communicative sounds while at the facility **102**. For ease of illustration and not necessarily by way of limitation the communicative sounds are referred to herein as “speech” and may include, but are not limited to, audible verbalizations, non-lexical utterances, and so forth. The non-lexical utterances may include sighs, gasps, grunts, clicks, pops, or other sounds generated by the user **116**. Other communicative sounds may include those generated by an action of the user **116**, such as fingersnaps, handclaps, foot stomps, and so forth.

The inventory management system **122** may include a speech interaction system **124** configured to access user context data **126** which is indicative of a user context to process user speech **128** in the audio data, and perform one or more actions. In some implementations, the one or more actions may include selecting one or more output devices **130** and processing user interface data to produce a user interface **132** for presentation to the user **116**. The speech interaction system **124** may include a speech recognition module, voice recognition module, a natural language understanding (NLU) module, a semantic analysis module, or other modules. The speech recognition module is configured to process audio of the user speech **128** and produce text or commands representative of the speech. The voice recognition module is configured to distinguish one user **116** from another user **116** based on the sounds emitted by the user **116**. The NLU module is configured to parse the information generated by the modules, such as from the speech recognition module, to generate information indicative of meaning of the user speech **128**. The NLU module may use Hidden Markov Models (HMMs), neural networks, or other techniques to generate speech from the audio. For example, neural networks may be used to pre-process the audio data and perform functions such as feature transformation, dimensionality reduction, and so forth, with the resulting output used for HMM-based recognition.

The user context data **126** may include information about one or more of: operation of the facility **102**, data about the items **104**, the user **116** as an anonymous individual, the user **116** as an identified individual, data external to the facility **102** such as weather or delivery route information, and so forth. For example, the user context data **126** may include position data indicative of a location of the one or more users **116** within the facility **102**. In another example, of the user

context data **126** may comprise user data such as user preferences which have been previously acquired and stored.

The user context data **126** may include data indicative of one or more items **104** within the facility **102** that the user **116** is interacting with. The interaction may include direct physical manipulation, such as picking up or touching the item **104**. The interaction may include looking at the item **104**, pointing, or gesturing at the item **104**. The interaction may include items **104** which are within a predetermined distance of the user **116**, such as within 2 meters.

In some implementations the user context data **126** may include information from other users **116** or groups of users **116** who have been determined to be similar in one or more respects. For example, the user context data **126** may include information such as user preferences as set by other users **116** with similar demographic data, who have similar access permissions, who were in about the same location within the facility **102**, and so forth. This information may be used in place of information which is not yet available about the user **116**, or may supplement information associated with the user **116**.

The user context data **126** may be used by the speech interaction system **124** to facilitate the functions thereof, such as recognizing words and determining meaning of the user speech **128**. The user context data **126** provides information which may provide cues to help disambiguate words, determine meaning, select particular actions, and so forth. For example, the user context data **126** may include a direction of a gaze of the user **116** which is indicative of a particular item **104** the user **116** is looking at. Using this information, the speech interaction system **124** may disambiguate the user speech **128** of “what is that” to determine that the user **116** is uncertain as to some attribute of the item **104**. For example, the user context may include a direction of a gaze of the user **116**, proximity to other users **116**, location of the user **116** within the facility **102**, and so forth. The user context may also include information which is specific to the user **116**. For example, the user context may include historical activity such as a list of items picked since entry to the facility **102** today.

The user context data **126** may be used by the speech interaction system **124** to modify statistical models, assign weights in neural networks, select speech recognition lexicons indicative of how words are pronounced, select speech recognition grammars indicative of an expected form or structure of spoken words, and so forth. In implementations where the speech interaction system **124** uses HMM-based recognition techniques, the user context data **126** may be used to configure one or more of the HMMs in the module. For example, the user context data **126** may be used to modify weightings in context dependency for phonemes, to assign weightings to possible transcriptions of the speech, and so forth. Continuing the example, the user context data **126** may be used to assess which ones of a set of possible transcriptions should be retained. Thus, where the user context data **126** indicates the user **116** has a gaze directed at a product in the category “ice cream”, the transcription of “ice cream” may be retained while transcriptions “I scream”, “eye scream”, “ice ream”, and “high screen” are disregarded.

A user device **134** may also be present within the facility **102**. The user device **134** may be provided by the user **116**, or may be provided by an operator of the facility **102** or another entity for use by the user **116**. The user device **134** may comprise a tablet computer, smartphone, notebook computer, data access terminal, wearable computing device,

and so forth. The user device **134** may be used to acquire audio data, process user interface data to present the user interface **132**, or provide other functions. For example, information from the user device **134** such as message data, call metadata, application state data, and so forth may be included in the user context data **126**.

The user **116** may benefit from various actions taken by the inventory management system **122** responsive to user speech **128**. For example, as the user **116(1)** moves about the facility **102** speaking to another person such as the user **116(2)** or another system, the speech interaction system **124** may process the user speech **128** and determine actions such as presenting information as output, moving inventory locations **114** or items **104**, and so forth. Instead of the user **116** explicitly querying inventory management system **122**, the system may instead determine and perform the actions unobtrusively. As a result, the experience of the user **116** in the facility **102** may be improved, efficiency of the facility's **102** operation may be improved, or other benefits may accrue.

FIG. **2** is a block diagram **200** illustrating additional details of the facility **102**, according to some implementations. The facility **102** may be connected to one or more networks **202**, which in turn connect to one or more servers **204**. The network **202** may include private networks, public networks such as the Internet, or a combination thereof. The network **202** may utilize wired technologies (e.g., wires, fiber optic cable, and so forth), wireless technologies (e.g., radio frequency, infrared, acoustic, optical, and so forth), or other connection technologies. The network **202** is representative of any type of communication network, including one or more of data networks or voice networks. The network **202** may be implemented using wired infrastructure (e.g., copper cable, fiber optic cable, and so forth), a wireless infrastructure (e.g., cellular, microwave, satellite), or other connection technologies.

The servers **204** may be configured to execute one or more modules or software applications associated with the inventory management system **122**. While the servers **204** are illustrated as being in a location outside of the facility **102**, in other implementations at least a portion of the servers **204** may be located at the facility **102**. The servers **204** are discussed in more detail below with regard to FIG. **3**.

The users **116**, the totes **118**, or other objects in the facility **102** may be equipped with one or more radio frequency (RF) tags **206**. The RF tags **206** are configured to emit an RF signal **208**. In one implementation, the RF tag **206** may be a radio frequency identification (RFID) tag configured to emit the RF signal **208** upon activation by an external signal. For example, the external signal may comprise a radio frequency signal or a magnetic field configured to energize or activate the RFID tag. In another implementation, the RF tag **206** may comprise a transmitter and a power source configured to power the transmitter. For example, the RF tag **206** may comprise a Bluetooth Low Energy (BLE) transmitter and battery. In other implementations, the tag may use other techniques to indicate presence. For example, an acoustic tag may be configured to generate an ultrasonic signal which is detected by corresponding acoustic receivers. In yet another implementation, the tag may be configured to emit an optical signal.

The inventory management system **122** may be configured to use the RF tags **206** for one or more of identification of the object, determining a position of the object, and so forth. For example, the users **116** may wear RF tags **206**, the totes **118** may have RF tags **206** affixed, and so forth which may be read and, based at least in part on signal strength,

used to determine identity and position. The tote **118** is configured to carry or otherwise transport one or more items **104**. For example, the tote **118** may include a basket, a cart, a bag, and so forth. The tote **118** is discussed in more detail below with regard to FIG. **4**.

Generally, the inventory management system **122** or other systems associated with the facility **102** may include any number and combination of input components, output components, and servers **204**.

The one or more sensors **120** may be arranged at one or more locations within the facility **102**. For example, the sensors **120** may be mounted on or within a floor, wall, or ceiling, at an inventory location **114**, on the tote **118**, in the user device **134**, may be carried or worn by the user **116**, and so forth.

The sensors **120** may include one or more imaging sensors **120(1)**. These imaging sensors **120(1)** may include cameras configured to acquire images of a scene. The imaging sensors **120(1)** are configured to detect light in one or more wavelengths including, but not limited to, terahertz, infrared, visible, ultraviolet, and so forth. The inventory management system **122** may use image data acquired by the imaging sensors **120(1)** during operation of the facility **102**. For example, the inventory management system **122** may identify items **104**, users **116**, totes **118**, and so forth based at least in part on their appearance within the image data.

One or more 3D sensors **120(2)** may also be included in the sensors **120**. The 3D sensors **120(2)** are configured to acquire spatial or three-dimensional data, such as depth information, about objects within a sensor field-of-view. The 3D sensors **120(2)** include range cameras, lidar systems, sonar systems, radar systems, structured light systems, stereo vision systems, optical interferometry systems, and so forth. The inventory management system **122** may use the three-dimensional data acquired to identify objects, determine a position of an object, and so forth. For example, the user context data **126** may include a position of the user **116** in three-dimensional space within the facility **102**.

One or more buttons **120(3)** may be configured to accept input from the user **116**. The buttons **120(3)** may comprise mechanical, capacitive, optical, or other mechanisms. For example, the buttons **120(3)** may comprise mechanical switches configured to accept an applied force from a touch of the user **116** to generate an input signal. The inventory management system **122** may use data from the buttons **120(3)** to receive information from the user **116**. For example, the tote **118** may be configured with a button **120(3)** configured such that the button **120(3)** may be activated by the user **116**.

The sensors **120** may include one or more touch sensors **120(4)**. The touch sensors **120(4)** may use resistive, capacitive, surface capacitance, projected capacitance, mutual capacitance, optical, Interpolating Force-Sensitive Resistance (IFSR), or other mechanisms to determine the position of a touch or near-touch. For example, the IFSR may comprise a material configured to change electrical resistance responsive to an applied force. The position of that change in electrical resistance within the material may indicate the position of the touch. The inventory management system **122** may use data from the touch sensors **120(4)** to receive information from the user **116**. For example, the touch sensor **120(4)** may be integrated with the tote **118** to provide a touchscreen with which the user **116** may select from a menu one or more particular items **104** for picking.

One or more microphones **120(5)** may be configured to acquire audio data indicative of sound present in the envi-

11

ronment. The sound may include the user speech **128** uttered by the user **116**. In some implementations, arrays of microphones **120(5)** may be used. These arrays may implement beamforming techniques to provide for directionality of gain. The speech interaction system **124** may use the one or more microphones **120(5)** to accept voice input from the users **116**, determine the position of one or more users **116** in the facility **102**, and so forth. For example, the microphones **120(5)** may use time difference of arrival (TDOA) techniques to localize a position of a sound.

One or more weight sensors **120(6)** are configured to measure the weight of a load, such as the item **104**, the user **116**, the tote **118**, and so forth. The weight sensors **120(6)** may be configured to measure the weight of the load at one or more of the inventory locations **114**, the tote **118**, or on the floor of the facility **102**. The weight sensors **120(6)** may include one or more sensing mechanisms to determine weight of a load. These sensing mechanisms may include piezoresistive devices, piezoelectric devices, capacitive devices, electromagnetic devices, optical devices, potentiometric devices, microelectromechanical devices, and so forth. The sensing mechanisms may operate as transducers which generate one or more signals based on an applied force, such as that of the load due to gravity. The inventory management system **122** may use the data acquired by the weight sensors **120(6)** to identify an object, determine a location of an object, maintain shipping records, and so forth.

The sensors **120** may include one or more light sensors **120(7)**. The light sensors **120(7)** may be configured to provide information associated with ambient lighting conditions such as a level of illumination. Information acquired by the light sensors **120(7)** may be used by the inventory management system **122** to adjust a level, intensity, or configuration of the output device **130**.

One more radio frequency identification (RFID) readers **120(8)**, near field communication (NFC) systems, and so forth may also be provided in the sensors **120**. For example the RFID readers **120(8)** may be configured to read the RF tags **206**. Information acquired by the RFID reader **120(8)** may be used by the inventory management system **122** to identify an object associated with the RF tag **206** such as the item **104**, the user **116**, the tote **118**, and so forth.

One or more RF receivers **120(9)** may also be provided. In some implementations the RF receivers **120(9)** may be part of transceiver assemblies. The RF receivers **120(9)** may be configured to acquire RF signals **208** associated with Wi-Fi, Bluetooth, ZigBee, 3G, 4G, LTE, or other wireless data transmission technologies. The RF receivers **120(9)** may provide information associated with data transmitted via radio frequencies, signal strength of RF signals **208**, and so forth. For example, information from the RF receivers **120(9)** may be used by the inventory management system **122** to determine a location of an RF source such as the user device **134**.

The sensors **120** may include one or more accelerometers **120(10)**, which may be worn or carried by the user **116**, mounted to the tote **118**, and so forth. The accelerometers **120(10)** may provide information such as the direction and magnitude of an imposed acceleration. Data such as rate of acceleration, determination of changes in direction, speed, and so forth may be determined using the accelerometers **120(10)**.

A gyroscope **120(11)** provides information indicative of rotation of an object affixed thereto. For example, the tote **118**, the user device **134**, or other objects may be equipped

12

with a gyroscope **120(11)** to provide user context data **126** indicative of a change in orientation.

A magnetometer **120(12)** may be used to determine a heading by measuring ambient magnetic fields, such as the terrestrial magnetic field. The magnetometer **120(12)** may be worn or carried by the user **116**, mounted to the tote **118**, the user device **134**, and so forth. For example, the magnetometer **120(12)** in the user device **134** as worn by the user **116(1)** may act as a compass and provide information indicative of which way the user **116(1)** is facing.

The sensors **120** may include other sensors **120(S)** as well. For example the other sensors **120(S)** may include proximity sensors, ultrasonic rangefinders, thermometers, barometric sensors, hygrometers, or biometric input devices including, but not limited to, fingerprint readers or palm scanners.

The facility **102** may include one or more access points **210** configured to establish one or more wireless networks. The access points **210** may use Wi-Fi, near field communication (NFC), Bluetooth, or other technologies to establish wireless communications between a device and the network **202**. The wireless networks allow the devices to communicate with one or more of the inventory management system **122**, the sensors **120**, the user devices **134**, the RF tag **206**, a communication device of the tote **118**, or other devices.

The output devices **130** may also be provided in the facility **102**. The output devices **130** are configured to generate signals which may be perceived by the user **116**.

Haptic output devices **130(1)** are configured to provide a signal which results in a tactile sensation of the user **116**. The haptic output devices **130(1)** may use one or more mechanisms such as electrical stimulation or mechanical displacement to provide the signal. For example, the haptic output devices **130(1)** may be configured to generate a modulated electrical signal which produces an apparent tactile sensation in one or more fingers of the user **116**. In another example, the haptic output devices **130(1)** may comprise piezoelectric or rotary motor devices configured to provide a vibration which may be felt by the user **116**.

One or more audio output devices **130(2)** are configured to provide acoustic output. The acoustic output includes one or more of infrasonic sound, audible sound, or ultrasonic sound. The audio output devices **130(2)** may use one or more mechanisms to generate the sound. These mechanisms may include, but are not limited to: voice coils, piezoelectric elements, magnetostrictive elements, or electrostatic elements, and so forth. For example, a piezoelectric buzzer or a speaker may be used to provide acoustic output.

The display devices **130(3)** may be configured to provide output which may be seen by the user **116**, or detected by a light-sensitive detector such as an imaging sensor **120(1)** or light sensor **120(7)**. The output may be monochrome or color. The display devices **130(3)** may be emissive, reflective, or both. An emissive display device **130(3)** is configured to emit light during operation. For example, a light emitting diode (LED) is an emissive visual display device **130(3)**. In comparison, a reflective display device **130(3)** relies on ambient light to present an image. For example, an electrophoretic display is a reflective display device **130(3)**. Backlights or front lights may be used to illuminate the reflective visual display device **130(3)** to provide visibility of the information in conditions where the ambient light levels are low.

Mechanisms of the display devices **130(3)** may include liquid crystal displays **130(3)(1)**, transparent organic light emitting diodes (LED) **130(3)(2)**, electrophoretic displays **130(3)(3)**, image projectors **130(3)(4)**, or other display mechanisms **130(3)(S)**. The other display mechanisms **130**

13

(3)(S) may include, but are not limited to, micro-electromechanical systems (MEMS), spatial light modulators, electroluminescent displays, quantum dot displays, liquid crystal on silicon (LCOS) displays, cholesteric displays, interferometric displays, and so forth. These mechanisms are configured to emit light, modulate incident light emitted from another source, or both.

The display devices 130(3) may be configured to present images. For example, the display devices 130(3) may comprise a pixel-addressable display. The image may comprise at least a two-dimensional array of pixels, or a vector representation of an at least two-dimensional image.

In some implementations, the display devices 130(3) may be configured to provide non-image data, such as text characters, colors, and so forth. For example, a segmented electrophoretic display, segmented LED, and so forth may be used to present information such as a SKU number. The display devices 130(3) may also be configurable to vary the color of the text, such as using multicolor LED segments.

In some implementations, display devices 130(3) may be configurable to provide image or non-image output. For example, an electrophoretic display 130(3)(3) with addressable pixels may be used to present images of text information, or all of the pixels may be set to a solid color to provide a colored panel.

The output devices 130 may include hardware processors, memory, and other elements configured to present a user interface 132. In one implementation, the display devices 130(3) may be arranged along the edges of inventory locations 114. The speech interaction system 124 may hear user speech 128 such as “where is that widget” and may be configured to present arrows on the display devices 130(3) which point to the inventory location 114 which contains the “widget” item 104.

Other output devices 130(T) may also be present. The other output devices 130(T) may include lights, scent/odor dispensers, document printers, three-dimensional printers or fabrication equipment, and so forth. For example, the other output devices 130(T) may include lights which are located on the inventory locations 114, the totes 118, and so forth. Continuing the example, the action of the speech interaction system 124 may comprise activating one or more lights on the tote 118 to navigate the user 116 through the facility 102 to a particular location.

The inventory management system 122 may generate the user interface data which is then used by the output device 130, the user devices 134, or both to present a user interface 132. The user interface 132 may be configured to stimulate one or more senses of the user 116. For example, the user interface 132 may comprise visual, audible, and haptic output.

FIG. 3 illustrates a block diagram 300 of the server 204. The server 204 may be physically present at the facility 102, may be accessible by the network 202, or a combination of both. The server 204 does not require end-user knowledge of the physical location and configuration of the system that delivers the services. Common expressions associated with the server 204 may include “on-demand computing,” “software as a service (SaaS),” “platform computing,” “network-accessible platform,” “cloud services,” “data centers” and so forth. Services provided by the server 204 may be distributed across one or more physical or virtual devices.

The server 204 may include one or more hardware processors 302 (processors) configured to execute one or more stored instructions. The processors 302 may comprise one or more cores. The server 204 may include one or more input/output (I/O) interface(s) 304 to allow the processor

14

302 or other portions of the server 204 to communicate with other devices. The I/O interfaces 304 may comprise integrated circuit (I2C), serial peripheral interface bus (SPI), Universal Serial Bus (USB) as promulgated by the USB Implementers Forum, RS-232, and so forth.

The I/O interface(s) 304 may couple to one or more I/O devices 306. The I/O devices 306 may include input devices such as one or more of a keyboard, mouse, scanner, the sensors 120, and so forth. The I/O devices 306 may also include output devices 130 such as one or more of a display, printer, audio speakers, and so forth. In some embodiments, the I/O devices 306 may be physically incorporated with the server 204 or may be externally placed.

The server 204 may also include one or more communication interfaces 308. The communication interfaces 308 are configured to provide communications between the server 204 and other devices, such as the sensors 120, the user devices 134, routers, the access points 210, and so forth. The communication interfaces 308 may include devices configured to couple to personal area networks (PANs), wired and wireless local area networks (LANs), wired and wireless wide area networks (WANs), and so forth. For example, the communication interfaces 308 may include devices compatible with Ethernet, Wi-Fi™, and so forth.

The server 204 may also include one or more busses or other internal communications hardware or software that allow for the transfer of data between the various modules and components of the server 204.

As shown in FIG. 3, the server 204 includes one or more memories 310. The memory 310 comprises one or more computer-readable storage media (“CRSM”). The CRSM may be any one or more of an electronic storage medium, a magnetic storage medium, an optical storage medium, a quantum storage medium, a mechanical computer storage medium, and so forth. The memory 310 provides storage of computer-readable instructions, data structures, program modules, and other data for the operation of the server 204. A few example functional modules are shown stored in the memory 310, although the same functionality may alternatively be implemented in hardware, firmware, or as a system on a chip (SOC).

The memory 310 may include at least one operating system (OS) module 312. The OS module 312 is configured to manage hardware resource devices such as the I/O interfaces 304, the I/O devices 306, the communication interfaces 308, and provide various services to applications or modules executing on the processors 302. The OS module 312 may implement a variant of the FreeBSD operating system as promulgated by the FreeBSD Project, other UNIX or UNIX-like variants, a variation of the Linux operating system as promulgated by Linus Torvalds, the Windows Server operating system from Microsoft Corporation of Redmond, Wash., and so forth.

Also stored in the memory 310 may be one or more of the following modules. These modules may be executed as foreground applications, background tasks, daemons, and so forth.

A communication module 314 may be configured to establish communications with one or more of the sensors 120, one or more of the user devices 134, other servers 204, or other devices. The communications may be authenticated, encrypted, and so forth.

The memory 310 may store an inventory management module 316. The inventory management module 316 is configured to provide the inventory functions as described herein with regard to the inventory management system 122. For example, the inventory management module 316 may

track items **104** between different inventory locations **114**, to and from the totes **118**, and so forth.

The inventory management module **316** may be configured to generate the user interface data which may be used by the output devices **130**, user devices **134**, and so forth to present the user interface **132**. The inventory management module **316** may include one or more of a data acquisition module **318**, a speech interaction module **320**, or a user interface module **322**.

The data acquisition module **318** is configured to acquire input from one or more of the sensors **120**. For example, the data acquisition module **318** may be configured to receive image data **328(1)** from the imaging sensors **120(1)** and audio data **328(2)** generated by the microphones **120(5)**. In some implementations, the data acquisition module **318** may analyze the sensor data **328** to generate additional information. For example, the data acquisition module **318** may use one or more techniques to localize or determine a position from which a sound emanates. These techniques may include beamforming using an array of a plurality of microphones **120(5)**, time difference of arrival (TDOA), analysis of a phase of sounds received by the array of microphones **120(5)**, and so forth.

The data acquisition module **318** may also be configured to acquire information from the user device **134**. In one implementation, the user device **134** may include an application module or operating system configured to provide information about messages sent or received by the user **116** of the user device **134**, application state such as whether a particular application is executing and if so what portion of the user interface **132** is being presented, and so forth. For example, the user device **134** may include an instant messaging application which is used to provide communication between users **116** of the facility **102** and other parties such as customers. The user device **134** in this example may be configured to provide to the data acquisition module **318** information indicative of messages sent or received by the user device **134**. The information acquired may be incorporated into the user context data **126** for use in processing the user speech **128** to perform one or more actions within the facility **102**.

The speech interaction system **124** may be implemented as the speech interaction module **320**. The speech interaction module **320** processes the sensor data **328** to determine intent, meaning, actors, or other aspects of the user speech **128**. Once determined, particular actions may be selected and then initiated. These actions may include, but are not limited to, storing data, retrieving data, activating one or more mechanisms within the facility **102**, directing other users **116** of the facility **102**, and so forth.

To facilitate the determination, the speech interaction module **320** may access the user context data **126**. The speech interaction module **320** processes user speech **128** which may not be specifically directed to the inventory management module **316**. Rather, the user speech **128** is incidental to the presence of the user **116** in the facility **102**. The user speech **128** may comprise the self-speak or muttering of the user **116** while alone, a conversation between the two users **116** in the facility **102**, a conversation between the user **116** in the facility **102** and another person elsewhere such as via telephone, and so forth. The natural language as expressed by the user **116** in the user speech **128** may be highly variable in terms of wordbase, lexicon, grammar, syntax, and so forth. For example, when the user **116** is specifically addressing a traditional speech recognition system the user **116** may speak with a particular syntax and grammar. In comparison, when the user **116** is speaking to

another person a looser or more complicated syntax and grammar may be used, such as that involving slang, idiom, allusion, personal references, and so forth.

To process the user speech **128**, the speech interaction module **320** may access the user context data **126** to determine a natural language understanding of the user speech **128**. The user context data **126** provides cues or information which may be used to determine wordbase, grammar, syntax, personal references, and so forth. For example, the user context data **126** may comprise information indicative of an item **104(1)** in the inventory location **114(1)** which the user **116(1)** is pointing at with an index finger. Based on this information, the speech interaction module **320** may process user speech **128** of "what's wrong with that?" to determine that the user **116** is inquiring about the status of the item **104(1)** in the inventory location **114(1)** and initiate an action which generates a user interface **132** which presents information about the status of the item **104(1)**.

The speech interaction module **320** may analyze the user speech **128** to determine whether the user speech **128** is directed to no one else (e.g., self-speech such as talking to oneself), to another user **116**, to an automated system, and so forth. The analysis may include assessment of attributes such as intonation, speaking rate, fluency, perplexity, and so forth to determine or suggest an intended recipient of the user speech **128**. For example, the intonation describes variation in spoken pitch as distinct from phonemes while the speaking rate indicates how fast the user **116** is generating communicative sounds. Perplexity may be indicative of relative complexity of the language models used to process the user speech **128**. For example, the perplexity may be indicative of, on average, how many different equally most probable words can follow any given word. Higher values of perplexity may be indicative of an increase in the information encoded in the user speech **128**.

One or more thresholds may be set for these attributes and used to distinguish self-speech, conversation with another user **116**, interaction with another automated system, and so forth. For example, user speech **128** which exhibits average speaking rate of above 120 words per minute, variation in intonation exceeding 10%, and high perplexity may be designated as user-to-user speech and may be analyzed accordingly using user context data **126** of the users **116** involved. In comparison, the user speech **128** which exhibits an average speaking rate of between 80 and 120 words, has intonation with variability below 10%, and a high perplexity may be deemed to be self-talk, such as the user **116** speaking to himself.

The user interface module **322** is configured to generate user interface data. The user interface data may be based at least in part on output from the speech interaction module **320**. For example, the speech interaction module **320** may retrieve information responsive to the user speech **128** and pass this information to the user interface module **322** for generation of user interface data and subsequent presentation of the user interface **132**. In some implementations, the user interface **132** may provide information associated with one or more of: contemporaneous user speech **128** or previous user speech **128**. For example, the user interface **132** may present visually on a display output device **130(3)** a listing of the last ten actions taken responsive to user speech **128**. The user interface **132** may also be configured to accept information from the user **116**, such as an indication as to whether the action taken was helpful. As described below, this input may be used to train the speech interaction module **320**.

17

As described above, the user interface data is configured to provide the user interface 132 by way of one or more output devices 130, user devices 134, or a combination thereof. The user interface 132 may include one or more of haptic, audible, or visual stimuli. For example, the user interface 132 may be audible such that synthesized speech is generated by the user interface module 322 and presented by the audio output devices 130(2).

In some implementations the speech interaction module 320 may be unable to resolve the user speech 128 to a single action. For example, the user speech 128 as interpreted with regard to the user context data 126 may have two or more equally valid actions. In some implementations a single action may be selected, or one or more of the actions may be taken. For example, the speech interaction module 320 may be unable to resolve whether the user 116 is having trouble finding the item 104 “widget” on the shelf, or whether the user 116 would like to change a quantity of the widgets for picking. The speech interaction module 320 may be configured to take multiple actions, such as providing navigational cues to the user 116 to find the item 104 and presenting in the user interface 132 controls which allow for changing the quantity of the widgets on the pick list.

Other modules 324 may also be present in the memory 310. For example, an object recognition module may be configured to use data from one or more of the sensors 120 to identify an object such as the item 104, the user 116, the tote 118, and so forth.

The memory 310 may also include a data store 326 to store information. The data store 326 may use a flat file, database, linked list, tree, executable code, script, or other data structure to store the information. In some implementations, the data store 326 or a portion of the data store 326 may be distributed across one or more other devices including other servers 204, network attached storage devices and so forth.

The data store 326 may also include sensor data 328. The sensor data 328 comprises information acquired from, or based on, the one or more sensors 120. For example, the sensor data 328 may comprise three-dimensional information about an object in the facility 102. As described above, the sensors 120 may include an imaging sensor 120(1) which is configured to acquire one or more images. These images may be stored as image data 328(1). The image data 328(1) may comprise information descriptive of a plurality of picture elements or pixels. The microphones 120(5) are configured to acquire audio data 328(2) which is representative of sounds at the facility 102. The audio data 328(2) may include user speech 128, movement sounds, equipment sounds, and so forth.

In some implementations, the sensor data 328 may be acquired at least in part by the user devices 134. For example, the data acquisition module 318 may receive image data 328(1) and audio data 328(2) from an imaging sensor 120(1) and a microphone 120(5) which are onboard a tablet computer of the user 116.

The data store 326 may also store the user context data 126. In some implementations, a portion of the user context data 126 may be retrieved from other data stores or devices. The user context data 126 is configured to provide information about one or more of: operation of the facility 102, data about the items 104, the user 116 as an anonymous individual, the user 116 as an identified individual, data external to the facility 102 such as weather or delivery route information, and so forth.

The user context data 126 may include physical layout data 330. The physical layout data 330 provides a mapping

18

of physical positions within the facility 102 of devices and objects such as the sensors 120, inventory locations 114, and so forth. For example, the physical layout data 330 may indicate the coordinates within the facility 102 of an inventory location 114, an RFID reader 120(8) close to that inventory location 114, and so forth. In some implementations, the inventory management module 316 may access the physical layout data 330 to determine a distance between two objects, such as the users 116(1) and 116(2).

Item data 332 may also be included in the user context data 126. The item data 332 comprises information associated with the items 104. The information may include one or more inventory locations 114, at which one or more of the items 104 are stored. The item data 332 may also include order data, SKU or other product identifier, price, quantity on hand, weight, expiration date, images of the item, detailed description information, ratings, ranking, and so forth. The inventory management module 316 may store information associated with inventory management functions in the item data 332. The speech interaction module 320 may be configured to retrieve information about the items 104 which are mentioned in, or relevant to, the user speech 128.

The user context data 126 may also include historical activity data 334, position data 336, proximity data 338, or gaze direction data 340. The historical activity data 334 provides information about the user's 116 interactions with the facility 102 and the objects therein, such as the items 104, the inventory locations 114, other users 116, the totes 118, and so forth. For example, the historical activity data 334 may indicate a time of entry to the facility 102, a route taken while in the facility 102, a list of items 104 which have been picked or placed, and so forth. The historical activity data 334 may be associated with the user 116 which has been identified or unidentified. In some implementations, the historical activity data 334 may be retrieved, such as in the circumstance of an identified user 116. For example, the user 116 may be identified upon entry to the facility 102. After identification, historical activity data 334 associated with a user account of the user 116 may be retrieved from the data store 326. The speech interaction module 320 may use the historical activity data 334 to ascertain what or whom the user 116 has interacted with in the facility 102.

The position data 336 comprises information indicative of the user's 116 position in the facility 102. The position data 336 may be expressed in absolute terms or relative terms. Absolute terms may comprise a latitude, longitude, and altitude with respect to a geodetic reference point. Relative terms may include a position of 25.4 m along an x-axis and 75.2 m along a y-axis as designated by a floor plan of the facility 102, 5.2 m from an inventory location 114 along the heading of 169°, and so forth. The position data 336 may be used by the speech interaction module 320 to determine user context data 126. For example, based on the position data 336 in the facility, and using the information about inventory locations 114 and the corresponding items 104 therein, the speech interaction module 320 may determine the user context data 126 includes the items 104 which are nearby.

The proximity data 338 provides information which is indicative of a spatial arrangement between the user 116 and another person. For example, the proximity data 338 may indicate that the user 116(1) is within a threshold distance of the user 116(2). In some implementations, the threshold distances may be specified for individual users 116. The threshold distances may be static or dynamically adjusted. For example, the threshold distance may be configured to 3 m from users 116 which are unrelated, and 1 m from the

users 116 which have a familial relationship. The speech interaction module 320 may use the proximity data 338 to process the user speech 128. For example, the user 116(1) may say the user speech 128 “what do you have to pick next?” Based at least in part on the proximity data 338 indicating that the user 116(2) is within a threshold distance and no other users 116 are within the threshold distance, the speech interaction module 320 may perform several actions. These actions may include accessing information indicative of the user’s 116(2) pick list and presenting a user interface 132 which displays the information about the next item 104 on the user device 134(2) associated with the user 116(2).

The gaze direction data 340 provides information indicative of a direction the user 116 is looking. In some implementations, the gaze direction data 340 may be determined by processing the image data 328(1) of one or more eyes of the user 116. Based at least in part on the position data 336 of the user 116 and the gaze direction data 340, it may be determined what object the user 116 is looking at. By using this information, the speech interaction module 320 may be able to disambiguate user speech 128. For example, the user 116(1) may be a supervisor who utters the user speech 128 “please get me another tote” while looking at the user 116(3) who is several meters away. The speech interaction module 320 may use the gaze direction data 340 to determine that the user 116(1) is addressing the user 116(3), and may perform the actions of providing a user interface 132 which is indicative of a closest tote 118.

In other implementations, other directional data associated with the user 116 may be determined. For example, the directional data may indicate overall orientation of the user’s 116 body, the direction of their head, and so forth.

The user context data 126 may also include user data 342. The user data 342 comprises information which is specific to an identified user 116. Once the user 116 has been identified, the user data 342 associated with that identity may be retrieved. The user data 342 may include demographic data 344, relationship data 346, user preferences 348, navigation path data 350, access permissions 352, or other information.

The demographic data 344 may comprise information indicative of the user’s 116 skills, address, contact information, age, and so forth. The relationship data 346 provides information indicative of an association between the user 116 and other users or people. For example, the relationship data 346 may indicate the users 116(1) and 116(2) are supervisor and subordinate, respectively. In another implementation, the relationship data 346 may indicate that the users 116(1) and 116(2) share a familial or marital relationship, such as being spouses or siblings. The relationship data 346 may be used by the inventory management module 316 in various ways. For example, the relationship data 346 may be accessed to prevent having related individuals crosscheck one another’s picking of a high-value item 104. The speech interaction module 320 may use the relationship data 346 to disambiguate, determine meaning, or otherwise process the user speech 128.

The user preferences 348 may include information such as preferred colors for visual display, preferences as to the timbre of a synthesized voice, preference for audible output over visual output, threshold distances for different actions, and so forth. For example, the user preference 348 may include threshold distance data indicating that the user 116(1) prefers a visual user interface 132 when other users 116 are within a 3 m threshold distance and an audible user interface 132 when the other users are beyond the 3 m threshold distance. The thresholds may be modified based on the physical layout data 330. For example, the threshold

distance may be within a “line of sight”, such as between users 116 which are within the same aisle 112 as compared to users 116 which are in adjacent aisles.

The navigation path data 350 provides information indicative of an actual, scheduled, or predicted path of the user 116 through the facility 102. For example, the historical navigation path data 350 may comprise a time series of waypoints of the user 116 within the facility 102 over a period of elapsed time. The waypoints may specify a location within the facility 102. The historical navigation path data 350 provides information showing where the user 116 has been within the facility 102.

Scheduled navigation path data 350 provides information as to what path the user 116 is being directed to follow through the facility 102. For example, the scheduled navigation path data 350 may be generated using a list of items 104 to be picked by the user 116. Based at least in part on the physical arrangement of the inventory locations 114 associated with those items 104 in the order as retrieved from the physical layout data 330, a picking path may be specified. In some implementations, the scheduled navigation path data 350 may also comprise a time series indicating estimated positions within the facility 102 at particular times. The routing of the scheduled navigation path data 350 may be configured to optimize one or more parameters of the user’s 116 travel through the facility 102. For example, the scheduled navigation path data 350 may be configured to minimize walking distance, reduce congestion at a particular inventory location 114 relative to other users 116, minimize use of other material handling equipment such as pallet jacks or forklifts, and so forth.

Predicted navigation path data 350 provides an estimation as to what path the user 116 will follow through the facility 102 prospectively. In some implementations, the predicted navigation path data 350 may also comprise a time series indicating estimated positions within the facility 102 at particular times. For example, the predicted navigation path data 350 may be based on the historical navigation path data 350 of one or more users 116. Continuing the example, the historical navigation path data 350 may indicate that a significant percentage of the users 116 that have gone past a particular aisle 112, enter that aisle 112 even when no items 104 within the aisle 112 are scheduled for picking. In this example, predicted navigation path data 350 for a particular user 116 may include the assumption that this particular user 116 will also enter that aisle 112.

In some implementations, certain portions of the facility 102 may be blanked out such that position data 336 is not collected, or may have navigation path data 350 associated with those portions redacted from the navigation path data 350. For example, movements within a break room or a restroom are not tracked.

The speech interaction module 320 may use the navigation path data 350 to establish where the user 116 may have been, or may be going next. For example, the user speech 128 may be the user 116(1) saying “my tote is so heavy” aloud. The speech interaction module 320 may determine the inventory location 114 which the user 116(1) is predicted to visit, and may dispatch another user 116(2) to meet the first user 116(1) and help with manipulating the heavily laden tote 118.

The user data 342 may also include access permissions 352 which specify what actions may be performed by a particular user 116 or group of users 116. For example, the user 116(1) may be assigned access permissions 352 commensurate with their responsibilities as a supervisor. These access permissions 352 may allow the supervisor user

116(1) to see detailed information about the tasks currently being performed by other users 116, allow for repositioning of inventory locations 114, and so forth.

In another example, the user 116(1) may have set access permissions 352 restricting dissemination of user data 342 to other users 116. The access permissions 352 may be specified in terms of which users 116 may retrieve user data 342, what user data 342 the particular user 116 chooses to allow others to access, or a combination thereof.

Individual users 116 or groups of users 116 may selectively provide user data 342 for use by the inventory management system 122, or may authorize collection of the user data 342 during use of the facility 102 or access to user data 342 obtained from other systems. For example, a user 116 may opt-in to collection of the user data 342 to receive enhanced services while using the facility 102.

The data store 326 may also store user interface data 354. The user interface data 354 may comprise commands, instructions, tags, markup language, images, color values, text, or other data. For example, the user interface data 354 may be expressed as hypertext markup language (HTML), cascading style sheets (CSS), JavaScript, and so forth. One or more output devices 130, user devices 134, or a combination thereof are configured to use the user interface data 354 to present the user interface 132 which may be perceived by the user 116. The user interface 132 may include one or more elements including visual, haptic, audible, olfactory, and so forth. For example, the user interface 132 may be a graphical user interface, audible user interface, haptic user interface, or a combination thereof.

The data store 326 may also include other data 356. For example, the other data 356 may include information such as configuration files, weather information, loading dock status, tote 118 availability, and so forth.

The server 204 may also include a power supply 358. The power supply 358 is configured to provide electrical power suitable for operating the components in the server 204.

FIG. 4 illustrates a block diagram 400 of the tote 118, according to some implementations. The tote 118 may include an RF tag 206. The RF tag 206 may be affixed to, integral with, or is otherwise associated with the tote 118. In some implementations, the tote 118 may have identifiers, tags, or other indicia thereupon. For example, a machine-readable optical code, such as a barcode, may be affixed to a side of the tote 118.

The tote 118 may include one or more hardware processors 402 (processors) configured to execute one or more stored instructions. The processors 402 may comprise one or more cores. The tote 118 may include one or more I/O interface(s) 404 to allow the processor 402 or other portions of the tote 118 to communicate with other devices. The I/O interfaces 404 may include I2C, SPI, USB, RS-232, and so forth.

The I/O interface(s) 404 may couple to one or more I/O devices 406. The I/O devices 406 may include one or more of the input devices such as the sensors 120. As described above, the sensors 120 may include imaging sensors 120(1), weight sensors 120(6), RFID readers 120(8), and so forth. The I/O devices 406 may also include output devices 130 such as haptic output devices 130(1), audio output devices 130(2), display devices 130(3), and so forth. For example, the tote 118 may include other display devices 130(T) such as lights which may be activated to provide information to the user 116. In some implementations, input and output devices may be combined. For example, a touchscreen display may incorporate a touch sensor 120(4) and a display

device 130(3). In some embodiments, the I/O devices 406 may be physically incorporated with the tote 118 or may be externally placed.

The tote 118 may also include one or more communication interfaces 408. The communication interfaces 408 are configured to provide communications between the tote 118 and other devices, such as other totes 118, user devices 134, routers, access points 210, the servers 204, and so forth. The communication interfaces 408 may include devices configured to couple to PANs, LANs, WANs, and so forth. For example, the communication interfaces 408 may include devices compatible with Ethernet, Wi-Fi, Bluetooth, Zig-Bee, and so forth.

The tote 118 may also include one or more busses or other internal communications hardware or software that allow for the transfer of data between the various modules and components of the tote 118.

As shown in FIG. 4, the tote 118 includes one or more memories 410. The memory 410 comprises one or more CRSM as described above. The memory 410 provides storage of computer readable instructions, data structures, program modules and other data for the operation of the tote 118. A few example functional modules are shown stored in the memory 410, although the same functionality may alternatively be implemented in hardware, firmware, or as a SOC.

The memory 410 may include at least one OS module 412. The OS module 412 is configured to manage hardware resource devices such as the I/O interfaces 404, the I/O devices 406, the communication interfaces 408, and provide various services to applications or modules executing on the processors 402. The OS module 412 may implement a variant of the FreeBSD operating system as promulgated by the FreeBSD Project, other UNIX or UNIX-like variants, a variation of the Linux operating system, such as Android as promulgated by Google, Inc. of Mountain View, Calif. Other OS modules 412 may be used, such as the Windows operating system from Microsoft Corporation of Redmond, Wash., the LynxOS from LynuxWorks of San Jose, Calif., and so forth.

Also stored in the memory 410 may be one or more of the following modules. These modules may be executed as foreground applications, background tasks, daemons, and so forth.

A communication module 414 may be configured to establish communications with one or more of the sensors 120, other totes 118, the user devices 134, the servers 204, or other devices. The communications may be authenticated, encrypted, and so forth.

The memory 410 may also store a tote item tracking module 416. The tote item tracking module 416 is configured to maintain a list of items 104, which are associated with the tote 118. For example, the tote item tracking module 416 may receive input from a user 116 by way of a touch screen display with which the user 116 may enter information indicative of the item 104 placed in the tote 118. In another example, the tote item tracking module 416 may receive input from the I/O devices 406, such as the weight sensor 120(6) and an RFID or NFC reader 120(8). The tote item tracking module 416 may send the list of items 104 to the inventory management system 122. The tote item tracking module 416 may also be configured to receive information from the inventory management system 122. For example, a list of items 104 to be picked may be presented within a user interface 132 on the display device 130(3) of the tote 118.

The memory 410 may include a display module 418. The display module 418 may be configured to present information, such as received from the one or more servers 204 or generated onboard the tote 118, using one or more of the output devices 130. For example, the display module 418 may be configured to receive user interface data 354 provided by the user interface module 322. By processing user interface data 354, the user interface 132 may be presented to the user 116 by way of the output devices 130 of the tote 118. For example, the user interface 132 may include haptic output from the haptic output devices 130(1), audio output from the audio output devices 130(2), images presented on the display output devices 130(3), activation of lights or other output devices 130(T), or a combination thereof. Continuing the example, the user interface data 354 may be generated by the speech interaction module 320 in response to processing of the user speech 128.

Other modules 420 may also be stored within the memory 410. In one implementation, a data handler module may be configured to generate sensor data 328. For example, an imaging sensor 120(1) onboard the tote 118 may acquire image data 328(1) and one or more microphones 120(5) onboard the tote 118 may acquire audio data 328(2). The sensor data 328, or information based thereon, may be provided to the data acquisition module 318.

The other modules 420 may also include a user authentication module which may be configured to receive input and authenticate or identify a particular user 116. For example, the user 116 may enter a personal identification number or may provide a fingerprint to the fingerprint reader to establish their identity.

The memory 410 may also include a data store 422 to store information. The data store 422 may use a flat file, database, linked list, tree, executable code, script, or other data structure to store the information. In some implementations, the data store 422 or a portion of the data store 422 may be distributed across one or more other devices including the servers 204, other totes 118, network attached storage devices and so forth.

The data store 422 may store a tote item identifier list 424. The tote item identifier list 424 may comprise data indicating one or more items 104 associated with the tote 118. For example, the tote item identifier list 424 may indicate the items 104 which are present in the tote 118. The tote item tracking module 416 may generate or otherwise maintain a tote item identifier list 424.

A unique identifier 426 may also be stored in the memory 410. In some implementations, the unique identifier 426 may be stored in rewritable memory, write-once-read-only memory, and so forth. For example, the unique identifier 426 may be burned into a one-time programmable non-volatile memory, such as a programmable read-only memory (PROM). In some implementations, the unique identifier 426 may be part of a communication interface 408. For example, the unique identifier 426 may comprise a media access control address associated with a Bluetooth interface. In some implementations, the user interface module 322 may use the unique identifier 426 to determine which tote 118 to generate the user interface 132 upon, or to determine a source for the sensor data 328.

The data store 422 may also store sensor data 328. The sensor data 328 may be acquired from the sensors 120 onboard the tote 118. The user interface data 354 received by the tote 118 may also be stored in the data store 422.

Other data 428 may also be stored within the data store 422. For example, tote configuration settings, user interface preferences, and so forth may also be stored.

The tote 118 may also include a power supply 430. The power supply 430 is configured to provide electrical power suitable for operating the components in the tote 118. The power supply 430 may comprise one or more of photovoltaic cells, batteries, wireless power receivers, fuel cells, capacitors, and so forth.

FIG. 5 illustrates a block diagram 500 of the user device 134, according to some implementations. The user device 134 may comprise a tablet computer, smartphone, notebook computer, data access terminal, wearable computing device, and so forth. The user device 134 may include one or more hardware processors 502 (processors) configured to execute one or more stored instructions. The processors 502 may comprise one or more cores. The user device 134 may include one or more I/O interface(s) 504 to allow the processor 502 or other portions of the user device 134 to communicate with other devices. The I/O interfaces 504 may include I2C, SPI, USB, RS-232, and so forth.

The I/O interface(s) 504 may couple to one or more I/O devices 506. The I/O devices 506 may include one or more of the input devices such as the sensors 120. As described above, the sensors 120 may include imaging sensors 120(1), 3D sensors 120(2), buttons 120(3), touch sensors 120(4), microphones 120(5), weight sensors 120(6), light sensors 120(7), RFID readers 120(8), accelerometers 120(10), gyroscopes 120(11), magnetometers 120(12), and so forth.

The I/O devices 506 may also include one or more of haptic output devices 130(1), audio output devices 130(2), display devices 130(3), and so forth. For example, the display devices 130(3) may be configured in a head-mounted display, such that the field of vision of the user 116 includes the display device 130(3) as well as one or more of the real-world objects in the facility 102. In some implementations, input and output devices may be combined. In some embodiments, the I/O devices 506 may be physically incorporated with the user device 134 or may be externally placed.

The user device 134 may also include one or more communication interfaces 508. The communication interfaces 508 are configured to provide communications between the user device 134 and other devices, such as other user devices 134, totes 118, routers, access points 210, the servers 204, and so forth. The communication interfaces 508 may include devices configured to couple to PANs, LANs, WANs, and so forth. For example, the communication interfaces 508 may include devices compatible with Wi-Fi, Bluetooth, ZigBee, and so forth.

The user device 134 may also include one or more busses or other internal communications hardware or software that allow for the transfer of data between the various modules and components of the user device 134.

As shown in FIG. 5, the user device 134 includes one or more memories 510. The memory 510 comprises one or more CRSM as described above. The memory 510 provides storage of computer readable instructions, data structures, program modules and other data for the operation of the user device 134. A few example functional modules are shown stored in the memory 510, although the same functionality may alternatively be implemented in hardware, firmware, or as a SOC.

The memory 510 may include at least one OS module 512. The OS module 512 is configured to manage hardware resource devices such as the I/O interfaces 504, the I/O devices 506, the communication interfaces 508, and provide various services to applications or modules executing on the processors 502. The OS module 512 may implement a variant of the FreeBSD operating system as promulgated by

the FreeBSD Project, other UNIX or UNIX-like variants, a variation of the Linux operating system, such as Android as promulgated by Google, Inc. of Mountain View, Calif. Other OS modules 512 may be used, such as the Windows operating system from Microsoft Corporation of Redmond, Wash., the LynxOS from LynuxWorks of San Jose, Calif., and so forth.

Also stored in the memory 510 may be one or more of the following modules. These modules may be executed as foreground applications, background tasks, daemons, and so forth.

A communication module 514 may be configured to establish communications with one or more of the sensors 120, user devices 134, the servers 204, or other devices. The communications may be authenticated, encrypted, and so forth.

The memory 510 may include a display module 516. The display module 516 may be configured to parse the user interface data 354 and present the user interface 132 using one or more of the output devices 130 which are onboard or communicatively coupled to the user device 134. For example, the display module 516 may include an HTML rendering engine configured to process the HTML in the user interface data 354. The display module 516 may also be configured to modify the user interface 132 based at least in part on sensor data 328 acquired by sensors 120 onboard or communicatively coupled to the user device 134. For example, the display module 516 may be configured to provide some object recognition and tracking, such that user interface elements presented by the display device 130(3) follow the motion of the tracked object, such as the user 116 which is in the field of view.

Other modules 518 may also be stored within the memory 510. In one implementation, a data handler module may be configured to generate sensor data 328. For example, the data handler module may be configured to acquire audio data 328(2) from one or more microphones 120(5) and provide that audio data 328(2) to the data acquisition module 318 of the server 204.

The other modules 518 may also include a facial detection module which may be configured to determine a presence of a human face within the image data 328(1) acquired by the onboard imaging sensor 120(1). The image of the face may then be sent to the server 204 for facial recognition to determine who is present in the image.

The memory 510 may also include a data store 520 to store information. The data store 520 may use a flat file, database, linked list, tree, executable code, script, or other data structure to store the information. In some implementations, the data store 520 or a portion of the data store 520 may be distributed across one or more other devices including the servers 204, other user devices 134, network attached storage devices and so forth.

The data store 520 may store sensor data 328 as acquired from the sensors 120 onboard the user device 134 or communicatively coupled thereto. For example, the sensor data 328 may include image data 328(1) acquired for an onboard imaging sensor 120(1) of the user device 134. The user interface data 354 received by the user device 134 may also be stored in the data store 520. Other data 522 may also be stored within the data store 520. For example, device configuration settings, user interface preferences, and so forth may also be stored.

The user device 134 may also include a power supply 524. The power supply 524 is configured to provide electrical power suitable for operating the components in the user device 134. The power supply 524 may comprise one or

more of photovoltaic cells, batteries, wireless power receivers, fuel cells, capacitors, and so forth.

FIG. 6 illustrates a side view 600 of an overhead imaging sensor 120(1) acquiring an image of the users 116, the tote 118, and other objects. In some implementations, the facility 102 may include one or more sensors 120 which are configured to acquire an image from an overhead vantage point or at other positions within the facility 102. The sensors 120 may include, but are not limited to, one or more of the imaging sensors 120(1), the 3D sensors 120(2), the microphones 120(5), the RFID readers 120(8), the RF receivers 120(9), and so forth. The sensor data 328 acquired by the sensors 120 may be used to generate the user context data 126.

In this illustration, one of the sensors 120 comprises an imaging sensor 120(1) which is configured to generate image data 328(1). A field of view 602 of the imaging sensor 120(1) depicted here includes the users 116(1) and 116(2) and the tote 118. The image data 328(1) may be provided to the inventory management system 122. For example, the inventory management module 316 executing on the server 204 may process the image data 328(1) to determine at least a portion of the user context data 126, such as the position data 336 for the users 116(1) and 116(2).

The one or more sensors 120 may also include an array of microphones 120(5) configured to generate audio data 328(2). The array of microphones 120(5) may be arranged in a one, two, or three-dimensional arrangement. For example, a three-dimensional arrangement may comprise microphones 120(5) arranged lengthwise along an aisle 112 at ceiling height and also lengthwise down the aisle on middle shelves of the inventory locations 114 for that aisle 112.

The array of microphones 120(5) may be used to localize or determine a position within the space of a particular sound. As described above, the localization may use time difference of arrival (TDOA), phase differences, amplitude comparison, beamforming, or other techniques to determine the position. The position may be determined along a single dimension, along two dimensions, or in three-dimensional space. For example, the position may be determined as a position lengthwise along the aisle 112 as a single dimension. Determination of the position data 336 may be based at least in part on the localization information provided by the array of microphones 120(5). For example, the localized position of the user speech 128 which is associated with the user 116(1) may be combined with position information generated from the image data 328(1) to provide the position data 336.

In some implementations, the data acquisition module 316 may be configured to generate user context data 126 using information from the plurality of sensors 120. For example, image data 328(1) may indicate that the user 116(1) is standing alone and speaking based on facial motions of the user 116(1) and an absence of other human faces. The audio data 328(2) may comprise the user speech 128. RF data from the RF receiver 120(9) as well as information provided by the user device 134 may be indicative of an active cell phone call. Furthermore, the information provided by the user device 134 may indicate that the user 116(1) is talking to the user 116(2) who is designated by the relationship data 346 as being a subordinate of the user 116(1). Based on this user context data 126, the speech interaction module 320 may be able to derive meaning and perform one or more actions responsive to the user speech 128 which comprises the half of the conversation uttered by the user 116(1). Thus, the speech interaction module 320 may be configured to initiate actions responsive to the user

speech 128 so long as at least one of the participants is present or otherwise in communication with the facility 102.

As described here, the tote 118 may also include one or more output devices 130, one or more sensors 120, and so forth. Similarly the user device 134 may include one or more output devices 130, one or more sensors 120, and so forth. The data acquisition module 318 may be configured to access the data provided by these sensors 120 while the user interface module 322 may be configured to generate user interface data 354 configured to provide the user interface 132 by way of the output devices 130.

FIG. 7 illustrates a user interface 700 for setting how the facility 102 interacts with the user 116, according to some implementations. One or more of the output devices 130 may present the user interface 700.

The user interface 700 may allow the user 116 to configure or otherwise set one or more of the user preferences 348. The user information section 702 provides information about the user 116 who is presently accessing the user interface 700 and their account information.

A user localization section 704 allows the user to view, set, or modify information such as a current location, preferred wordbase, and so forth. For example, as shown here the user 116 has a current location of North America indicative of the facility 102 being located within the United States. The preferred word base, in comparison, is that of the United Kingdom. Continuing the example, the user 116 may be a native of the United Kingdom. The wordbase may comprise words and their pronunciations, associated meanings, grammar, and so forth which may have particular regional, cultural, or other significance. For example, a United Kingdom wordbase may indicate that the word “torch” is equivalent to the American word for “flashlight”. The wordbase may be determined manually such as by input from the user 116, or may be automatically determined based on analysis of the user speech 128 by the user 116. In some implementations the word base may comprise a data structure which specifies interconnections or relationships between one or more lexicons, grammars, semantic data, and so forth.

The user specific keywords section 706 allows the user to view, set, or modify particular key words or phrases. Key-words or phrases may be associated with particular actions, intents, or states of the user 116. The association of particular keywords to the user 116, groups of users 116, categories of users 116, and so forth, may be automatically determined by the speech interaction module 320. For example, a grunt noise may be associated with frustration of a particular user 116. In other implementations, the user interface 700 may allow the user 116 to specify that the grunt conveys this particular meaning. Responsive to this meaning, the speech interaction module 320 may be configured to provide additional information associated with a task assigned to the user 116 upon processing of the user speech 128 which includes a grunt.

A user specific settings section 708 allows the user to specify other aspects of operation of the speech interaction module 320 or the inventory management module 316. For example, the user 116 may give permission for the speech interaction module 320 to continuously monitor the audio data 328(2) at all times. The user 116 may provide specific limitations such as including or excluding telephone calls, including or excluding conversations with children, and so forth. Other settings may include enabling the user gaze tracking, enabling operation with an application on the user device 134, and so forth.

The user specific settings section 708 may also allow the user 116 to specify whether information provided in the user interface 132 is to be personalized. Personalized information may include the user data 342. For example, a non-personalized user interface 132 may present information such as a SKU and cost of an item 104. In comparison, a personalized user interface 132 may present information such as a last quantity purchased by that user 116.

The user specific settings section 708 may allow the user 116 to specify whether audible output is to be provided, and if so whether to limit that audible output based on proximity data 338 or other parameters. Similarly, the user 116 may specify whether visual output is to be provided, and if so what limits with regard to proximity data 338 are to be applied.

The user 116 may also specify user specific settings section 708 which enables the speech interaction module 322 to engage in an interactive conversation. For example, the user speech 128 may result in retrieval of information about an item 104. The speech interaction module 320 may operate in conjunction with the user interface module 322 to generate the user interface 132 which presents synthesized speech or a visual prompt asking the user 116 for permission to present the retrieved information when this option is enabled. In another example, the speech interaction module 320 may operate in conjunction with the user interface module 322 to ask the user 116 questions or to provide additional information to disambiguate or confirm an action. When not enabled, the system may generate and present the user interface 132 automatically.

The user specific settings section 708 may also enable the speech interaction module 320 to learn from the interactions with the user 116. For example, one or more machine learning techniques may be employed to process speech and determine which actions to provide for particular users 116, all users 116, or a combination thereof. For example, the speech interaction module 320 may process user speech 128 using a neural network to recognize text and determine an action associated therewith.

Feedback may be provided to the machine learning system by monitoring further user speech 128, user context data 126, or a combination thereof. For example, subsequent user speech 128 such as “thank you”, “great,” or “that’s helpful” may provide data which reinforces that the action determined by the speech interaction module 320 was appropriate with respect to the user speech 128. In comparison, a user speech 128 such as “that’s wrong” or “nope” may provide data which deprecates the action.

Similarly, subsequent actions by the user 116 following an action by the inventory management system 122 may provide feedback to the speech interaction module 320. For example, the action of moving the item 104(1) close to the user 116, who then subsequently picks up the item 104(1), may reinforce that the user 116 wanted the item 104(1) and that the action of moving the item 104(1) was useful. In another example, the action of the user 116 moving toward the item 104, or the inventory location 114 associated with the item 104, provides feedback which may be analyzed. Similarly, user input from the user interface 132 may also be used to provide feedback to the speech interaction module 320.

Other sections which are not depicted here may be provided in the user interface 700. For example, an access permission section may allow the user 116 to specify one or more of the access permissions 352.

FIG. 8 illustrates a scenario 800 of a single user 116 uttering user speech 128 and the system 100 responding to

this, according to some implementations. As described above, the facility 102 may include one or more sensors 120 such as microphones 120(5) and one or more output devices 130 such as audio output devices 130(2). In this illustration, a single user 116 is shown in a portion of the facility 102 which includes an inventory location 114 holding various items 104. While working by himself, the user 116 utters user speech 128 such as “widgets, widgets, widgets, (hmmm)”. The speech interaction module 320, based on the user context data 126, may interpret the user speech 128 as calling for the action of providing information to the user 116 about the location of the items 104 known as widgets.

As illustrated here, the audio output device 130(2) provides a user interface 132(1) comprising synthesized audio saying “Allow me to help you. The widgets on your pick list are located to your left on rack A shelf 2.” Another user interface 132(2) comprising text presented visually reading “Widgets are on rack A shelf 2” and presented on a display output device 130(3) may be presented instead of, or in addition to, the audible user interface 132(1). In some implementations the user 116 may be advised that the speech interaction system 124 has determined a suggested action or is ready to perform an action by a chime, haptic output, change in illumination of a light, and so forth. For example, the speakers on the tote 118 may chime while lights on the tote 118 may blink. The action may be performed, or may be held pending approval or assent by the user 116 to perform the action.

As illustrated in this example, the user 116 has not explicitly engaged the system for operation. Indeed, the utterance of the user speech 128 may be accidental or unintentional and may lack a defined grammar or syntax. However, based on the user context data 126 which includes information about the pick list of items associated with the user 116, as well as the item data 332, the physical layout data 330, the user position data 336, and so forth, the speech interaction module 320 is able to provide useful information to the user 116 to facilitate performance of the user’s 116 tasks.

FIG. 9 illustrates a scenario 900 in which the system 100 responds to user speech 128 between two users 116(1) and 116(2) of the facility 102, according to some implementations. As described above, the facility 102 may include one or more sensors 120 such as microphones 120(5) and one or more output devices 130 such as audio output devices 130(2). In this illustration, two users 116(1) and 116(2) are shown in a portion of the facility 102 which includes an inventory location 114 holding various items 104. While working together, the first user 116(1) says aloud to the second user 116(2) “Bob, it seems like these widgets have gotten more expensive.” In terms of wording and inflection, this user speech 128 is declarative. However, the speech interaction module 320, based on the user context data 126, retrieves information about cost changes associated with the item 104 which is named “widget” and determines that an appropriate action is to present this information. The speech interaction module 320 may then provide this information to the user interface module 322 which generates user interface data 354 suitable for providing one or more of an audible user interface 132(1) or a visual user interface 132(2). The user 116(1) may provide an affirmation that the user interface 132 providing information was of use by issuing further user speech 128 such as “thank you”.

In another implementation not illustrated here, the user 116(1) may be carrying on the conversation with a person or system not present at that position within the facility 102 or not in the facility 102 at all. For example, the user 116(1)

may be conversing with the user 116(2) via a two-way radio, cell phone, or other audio communication device. In another example, the user 116(1) may be providing verbal input to another system, such as providing instructions or a query to an intelligent personal assistant application executing at least in part on the user device 134.

Illustrative Processes

FIG. 10 depicts a flow diagram 1000 of a process for performing actions based on speech between users 116 of the facility 102, according to some implementations. In some implementations, the process may be performed at least in part by the inventory management module 316.

Block 1002 acquires images within the facility 102. For example, the imaging sensors 120(1) may generate image data 328(1).

Block 1004 acquires audio within the facility 102. For example, the microphones 120(5) may generate the audio data 328(2).

Block 1006 identifies a first user 116(1) with the acquired images. For example, the image data 328(1) may be processed using one or more techniques to determine identity of the first user 116(1). These techniques may include one or more of facial recognition, clothing recognition, gait recognition, and so forth. In other implementations, other techniques may be used instead of or in addition to identification by way of the image data 328(1). For example, a user 116 may be identified by way of an RFID tag, manual entry of credentials at an input device, and so forth.

Facial recognition may include analyzing facial characteristics which are indicative of one or more facial features in an image, three-dimensional data, or both. The facial features include measurements of, or comparisons between, facial fiducials or ordinal points. The facial features may include eyes, mouth, lips, nose, chin, ears, face width, skin texture, three-dimensional shape of the face, presence of eyeglasses, and so forth. In some implementations the facial characteristics may include facial metrics. The facial metrics indicate various ratios of relative sizes and spacing of the facial features. For example, the facial metrics may include ratio of interpupillary distance to facial width, ratio of eye width to nose width, and so forth. In some implementations the facial characteristics may comprise a set of eigenvectors by using principal component analysis (PCA) on a set of images. These eigenvectors as descriptive of a human face may be known as “eigenfaces” or “eigenimages”. In one implementation the facial recognition described in this disclosure may be performed at least in part using one or more tools available in the OpenCV library as developed by Intel Corporation of Santa Clara, Calif., Willow Garage of Menlo Park, Calif., and Itseez of Nizhny Novgorod, Russia, with information available at www.opencv.org. In other implementations, other techniques may be used to recognize faces. Previously stored registration data may associate particular facial characteristics with a particular identity, such as represented by a user account. For example, the particular pattern of eigenvectors in the image may be sought in the previously stored data, and matches within a threshold tolerance may be determined to indicate identity. The facial characteristics may be used to identify the user 116, or to distinguish one user 116 from another.

Clothing recognition analyzes images to determine what articles of clothing, ornamentation, and so forth the user 116 is wearing or carrying. Skin and hair detection algorithms may be used to classify portions of the image which are associated with the user’s 116 skin or hair. Items which are not skin and hair may be classified into various types of articles of clothing such as shirts, hats, pants, bags, and so

31

forth. The articles of clothing may be classified according to function, position, manufacturer, and so forth. Classification may be based on clothing color, texture, shape, position on the user **116**, and so forth. For example, classification may designate an article of clothing worn on the torso as a “blouse” while color or pattern information may be used to determine a particular designer or manufacturer. The determination of the article of clothing may use a comparison of information from the images with previously stored data. Continuing the example, the pattern of the blouse may have been previously stored along with information indicative of the designer or manufacturer.

In some implementations, identification of the user **116** may be based on the particular combination of classified articles of clothing. For example, the user **116(2)** may be distinguished from the user **116(3)** based at least in part on the user **116(2)** wearing a hat and a red shirt while the user **116(3)** is not wearing a hat and is wearing a blue shirt. The clothing may be used to identify the user **116**, or distinguish one user **116** from another.

Gait recognition analyzes images, three-dimensional data, or both to assess how a user **116** moves over time. Gait comprises a recognizable pattern of movement of the user’s **116** body which is affected by height, age, and other factors. Gait recognition may analyze the relative position and motion of limbs of the user **116**. Limbs may include one or more arms, legs, and in some implementations, head. In one implementation edge detection techniques may be used to extract a position of one or more limbs of the user **116** in the series of images. For example, a main leg angle of a user’s **116** leg may be determined, and based on the measurement of this main leg angle over time and from different points-of-view, a three-dimensional model of the leg motion may be generated. The change in position over time of the limbs may be determined and compared with previously stored information to determine an identity of the user **116**, or distinguish one user **116** from another.

In some implementations, identity may be based on a combination of these or other recognition techniques. For example, the user **116** may be identified based on clothing recognition, gait recognition, and facial recognition. The different recognition techniques may be used in different situations, or in succession. For example, clothing recognition and gait recognition may be used at greater distances between the user **116** and the imaging sensors **120(1)** or when the user’s **116** face is obscured from view of the imaging sensor **120(1)**. In comparison, as the user **116** approaches the imaging sensor **120(1)** and their face is visible, facial recognition may be used. Once identified, such as by way of facial recognition, one or more of gait recognition or clothing recognition may be used to track the user **116** within the facility **102**.

In other implementations other techniques for facial recognition, clothing recognition, gait recognition, and so forth may be employed. For example, facial recognition may use iris detection and recognition, clothing recognition may use embedded RF tags, gait recognition may utilize data from weight sensors **120(6)**, and so forth.

Block **1008** processes the audio to detect user speech **128** spoken by the first user **116(1)** and directed to a second user **116(2)** of the facility **102**. For example, the audio data **328(2)** may be processed to isolate the sound of the first user’s **116(1)** voice and to remove background noises such as that of machinery in the facility **102**.

In some implementations, various techniques may be used to determine which user **116** is speaking within the facility **102**. For example, should the facility **102** be busy, multiple

32

users **116** may be speaking simultaneously. These techniques may include TDOA, beamforming, use of data from imaging sensors **120(1)**, weight sensors **120(6)**, and so forth. By determining which user **116** is speaking, a speech interaction module **320** may be able to access the user context data **126** associated with that user **116**.

The determination as to whether the first user **116(1)** is speaking or not may be based at least in part on one or more of: facial motions of the first user **116** in the acquired images, voice recognition, or sound source localization using the audio from the plurality of microphones **120(5)**. For example, timestamps of the sensor data **328** may be compared to determine if the image data **328(1)** for an image frame in which the user’s **116** mouth is open corresponds to a point in time at which audio data **328(2)** has been acquired from a microphone **120(5)** which is proximate to the position of the user **116** as specified in the position data **336**. By detecting the motion of the user’s face **116**, and the presence of corresponding audio, it may be determined that the user **116** is emitting a sound. In another example, voice recognition may be used to distinguish the speech of one user **116** from another. For example, the voice recognition may use pitch, pace, cadence, and so forth to distinguish one user **116** from another. In yet another example, using one or more beamforming techniques, or by sampling data from a plurality of microphones **120(5)** some of which are close to the user **116**, some of which are far away, the system may determine that a particular user **116** is speaking.

Block **1010** determines a first user context associated with the first user **116(1)**. For example, the user context data **126(1)** for the identified user **116(1)** may be retrieved from the data store **326**. The first user context data **126** may be indicative of one or more of: the historical activity data **334** of the first user **116(1)**, the gaze direction data **340** of the first user **116(1)**, position data **336** of the first user **116(1)** within the facility **102**, proximity data **338** of the first user **116(1)** to the second user **116(2)**, demographic data **344** of the first user **116(1)**, user preferences **348** of the first user **116(1)**, or other information.

In some implementations, the user context data **126** may include information indicative of the identity of other users **116**, such as those who are proximate to the first user **116(1)**. By knowing to whom the first user **116(1)** is speaking, the speech interaction module **320** may be better able to process the user speech **128** as described next.

Furthermore, in some implementations a second user context data **126(2)** associated with a second user **116(2)** may be determined. The second user **116(2)** may be identified as described as above. For example, the second user context data **126(2)** for the second user **116(2)** may be retrieved from the data store **326**. The second user context data **126(2)** may include the identity of the first user **116(1)**. Stated another way, each of the users **116(1)** and **116(2)** may have their own respective user context data **126** which is indicative of the identity of the other user **116**.

Block **1012** processes, using the first user context **126(1)**, the user speech **128** to determine one or more key words or phrases. The key words or phrases may be designated by a system administrator, designated by the user **116**, automatically determined, or a combination thereof. For example, the key words may be the name of an item **104** as stored in the facility **102**, particular adjectives such as “expensive”, particular nouns such as “tote”, and so forth. Continuing the example, as described above with regard to FIG. **9**, the keywords may be “widget”, “cost”, “expensive” and “last time”.

33

In other implementations where the second user context **126(2)** is available, the processing of the user speech **128** may be further based at least in part on the second user context **126(2)**. For example, by knowing the participants in the conversation, the speech interaction module **320** may be better able to determine actions.

Block **1014** retrieves data responsive to the one or more key words. For example, the speech interaction module **320** may generate a query for cost changes associated with the item **104** named “widget”.

Block **1016** presents output representative of the data with the output device **130**. Continuing the example, the output device **130** may comprise an audio output device **130(2)**. The output may comprise an audible user interface **132** in which the inventory management module **316** generates user interface data **354** configured to present the audible user interface **132(1)**, the visual user interface **132(2)**, or both responsive to the user speech **128**.

As described above, in some implementations the system **100** may be responsive to feedback provided by the user **116**. For example, after receiving the information about the widget, the user **116(1)** may find the information of use and respond by saying “thank you”. The second user speech **128(2)** as uttered by the first user **116(1)** may be deemed indicative of acceptance of the data provided to the first user **116(1)**. Information indicative of the acceptance may then be stored as historical activity data **334** of the first user **116(1)**. This information may then be used to modify operation of the speech interaction module **320**, such as by changing weights of one or more elements within a neural network.

In some implementations, the tote **118** may include one or more output devices **130**. For example, the tote **118** may include a display output device **130(3)**. This output device **130** onboard or coupled to the tote **118** may be used to present a visual user interface **132(2)**.

FIG. **11** depicts a flow diagram **1100** of another process for performing actions based on user speech **128** between users **116** of the facility **102**, according to some implementations.

Block **1102** determines a user context associated with a first user **116(1)** of the facility **102**. As described above, the user context data **126** may comprise data indicative of one or more of: the historical activity data **334** of the first user **116(1)**, position data **336** of the first user **116(1)** within the facility **102**, proximity data **338** of the first user **116(1)**, gaze direction data **340** of the first user **116(1)** to other users such as a second user **116(2)**, and so forth.

In some implementations, the user context data **126** may include relationship data **346**. The relationship data **346** may be indicative of a relationship between the first user **116(1)** and the second user **116(2)**. The relationship data **346** may be determined based on one or more of: the first user **116(1)** and the second user **116(2)** entering the facility **102** together, the first user **116(1)** and the second user **116(2)** moving together within the facility **102**, or previously stored information indicative of a relationship. For example, facility entry data may comprise log data or other information which indicates the first user **116(1)** and the second user **116(2)** entered the facility **102** contemporaneously and within a threshold distance of one another. In another example, the user context data **126** may indicate that the first user **116(1)** and the second user **116(2)** have a relationship because they have been walking proximate to one another for a predetermined length of time within the facility **102**.

The user context data **126** may be received, at least in part, from the user device **134**. For example, the data acquisition module **318** may receive from the user device **134** one or

34

more of information exchanged using a short message service (SMS), information exchanged using an instant messaging service, call metadata associated with an incoming or outgoing telephone or video call, application state data indicative of an application executing on the user device **134**, or other information. For example, the user context data **126** may include text from an instant message “Get more widgets”, while the user speech **128** may be the first user **116(1)** muttering “well, where are they?” By using the text of the instant message, the speech interaction module **320** is able to process the user speech **128** and determine that the first user **116(1)** would benefit from information indicative of the inventory locations **114** of the widgets.

In another example, the call metadata may comprise information such as a telephone number or other network address of the other party(s) with whom communication has been attempted or is established. Other call metadata may include a duration of the call, indication of who originated the call, and so forth.

In yet another example, the application state data may include information such as a browser history, look up history, and so forth. Continuing this example, the first user **116(1)** may be looking up technical details about widgets. Based on this information and the user speech **128**, the speech interaction module **320** may provide the information as to the whereabouts of the widgets.

The user context data **126** may include information indicative of a wordbase in use by the user **116**. The wordbase may comprise words, idioms, and phrasing information associated with how the user **116** uses language. Users **116** from different regions or backgrounds may use different wordbases. The speech interaction module **320** may be configured to determine a wordbase of the user **116** based on one or more of word selection of the user **116**, language of the user speech **128**, cadence of the user speech **128**, manual information designating a wordbase of the user **116**, and so forth. For example, the speech interaction module **320** may use demographic data **344** about the user **116(1)** to determine that the user **116(1)** is from the United Kingdom. As a result, the user speech **128** may be processed using a United Kingdom wordbase, enabling the speech interaction module **320** to correctly determine that the user **116(1)** is referring to a flashlight and not a combustible object in the phrase “where is the torch?”

As described above, in some implementations the process may include determining an identity of the first user **116(1)**. Once identified, user data **342** associated with the identity may be accessed. For example, facial recognition, speech recognition, receipt of an RFID tag, and so forth may be used to associate the first user **116(1)** with a particular user account and the corresponding user data **342** associated therewith.

Block **1104** acquires audio of the user speech **128** or other sounds uttered or otherwise generated by the first user **116(1)**. For example, the microphones **120(5)** may generate audio data **328(2)**. As described above, the user speech **128** may be directed to a second user **116(2)**. The second user **116(2)** may be an identified user of the facility **102**, a guest of the first user **116(1)**, or a person whom the first user **116(1)** is speaking to by way of a two-way radio, telephone, and so forth.

Block **1106** uses the first user context data **126(1)** to process the user speech **128**. The processing may include speech recognition, semantic analysis, and so forth. The processing results in a determination of one or more actions which the system **100** may perform to facilitate the user **116**.

35

The actions may include one or more of doing nothing, storing information, retrieving information, presenting information, dispatching assistance from another user **116**, moving inventory **104**, moving inventory locations **114**, and so forth.

Block **1108** performs the one or more actions. For example, the one or more actions may include moving one or more inventory locations **114** from a first position in the facility **102** to a second position in the facility **102**. As described above, the inventory locations **114** may be able to be repositioned within the facility **102**, such as moving from one aisle **112** to another. In another example, the one or more actions may include moving one or more items **104** from a first inventory location **114** to a second inventory location **114**, or from a first position of the facility **102** to a second position. For example, a conveyor, robot, crane, and so forth may be used to move the widget from a high storage rack to a lower level which is accessible to the first user **116(1)**.

Other actions may include modifying how the first user **116(1)** is directed through the facility **102**. A block or module may be configured to access a first proposed route for the first user **116(1)** through the facility **102**. For example, the navigation path data **350** data for the first user **116(1)** may be retrieved. Another block may then modify the first proposed route to generate a second proposed route. The modification may include adding or removing waypoints, changing a waypoint sequence, and so forth. For example, the modified path may include an addition of a waypoint for the inventory location **114** associated with the item **104** “widget” which the first user **116(1)** had mentioned in his user speech **128**.

The actions may further include providing navigational guidance in the facility **102** to the first user **116(1)** based on the second proposed route. The navigational guidance may comprise prompts provided by way of an output device which serves to direct the user **116(1)** to a particular location in the facility **102**. Continuing the example, the first user **116(1)** may be provided with audible prompts stating which direction to turn or visual prompts such as arrows or lines indicating a direction of travel corresponding to the second proposed route. In other implementations navigational guidance may comprise visual projections being projected within the facility **102**, illumination of particular lights within the facility **102**, and so forth.

As described above, the actions may include presenting information by way of the user interface **132**. In some implementations, the user interface **132** and how the user interface **132** is presented may be specific to the user **116**. For example, the user data **342** may include user preferences **348** which specify particular situations within which the visual user interfaces **132(2)** are to be used and other situations within which audible user interfaces **132(1)** are to be used. Based at least in part on the user data **342**, user interface data **354** may be generated. Also based at least in part on the user data **342**, a particular output device **130** may be selected. Continuing the example above, the first user **116(1)** may specify a user preference **348** that when another person is within 2 m of the first user **116(1)**, the visual interface is to be used and presented on the display output device **130(3)** of the user device **134** associated with the first user **116(1)**. Given that another person is present within the 2 m threshold distance, the user interface data **354** may be configured to provide a visual user interface **132(2)**. The user interface data **354** may be provided to the user device **134** which has been selected for presentation.

In some situations, the determined action may be modified based on the presence of a third person. For example,

36

the first user **116(1)** and the second user **116(2)** may be conversing with one another when a third user **116(3)** approaches. The block may determine a third person, such as the user **116(3)**, is within a threshold distance of the first user **116(1)**. For example, the threshold distance may be 3 m and the third user **116(3)** may be 2.5 m away from the first user **116(1)**, and thus within the threshold distance. User interface data **354** may be generated and may be sent to the output device **130** which is associated with the first user **116(1)**. The output device **130** may be associated in some implementations by being within the presence of or in the possession of the first user **116(1)**. For example, the user interface data **354** may be sent to the user device **134** which is being carried by the first user **116(1)**. In another example, the user interface data **354** may be sent to the tote **118** which is being used by the first user **116(1)**.

The user speech **128** may comprise an inquiry or may result in an inquiry. The one or more actions may include presentation of information responsive to the inquiry. The information presented in a user interface **132** may be strongly suggestive of the nature of the inquiry, or provide other information about the user context. It may be advantageous to safeguard information about the user context data **126** by obfuscating information presented in the user interface **132**. This obfuscation may include presenting information in the user interface **132** which is responsive to the original inquiry, as well as other information which has been previously inquired about the other users **116**. Thus, the user interface **132** provides the information requested by the user **116**, as well as information which is not a request by the user **116**. To another party observing the user interface **132**, it may not be apparent what information is specifically associated with the user **116** which made the initial inquiry.

One or more blocks execute a first query to retrieve first information associated with the first inquiry. The first query may be responsive to user speech **128** uttered by the first user **116(1)**. A second query is executed to retrieve second information from a previous inquiry made by a different user **116**, such as from user **116(5)**. In some implementations the second query may be selected from actions corresponding to user speech **128** previously uttered in other similar circumstances, while in the same area of the facility **102**, and so forth. The user interface data **354** is generated using the first information and the second information.

An output device **130** accessible to the first user **116(1)** and other users of the facility **102** is selected. For example, the output device **130** may comprise an overhead speaker or a publicly visible display device **130(3)**. The user interface data **354** is processed to provide a user interface **132** by way of the selected output device **130**. For example, the user interface **132** may provide the information responsive to the user speech **128** as well as the other information.

Some information which may be presented by the user interface **132** may be deemed sensitive, such that public presentation is discouraged or impermissible. The actions performed by the inventory management module **316** may include one or more of presenting information designated as personally identifiable information, presenting information designated as sensitive, or presenting information about another user **116**, and so forth. For example, the personally identifiable information may include a residential address of the user **116**, information about a customer of the facility **102**, and so forth.

In such a situation, an output device **130** associated with the first user **116(1)** may be selected. The association may be determined based at least in part on proximity, user preferences **348**, and so forth. For example, the user preference

348 may specify that sensitive information is to be presented by way of the user device 134 of the first user 116(1), instead of a display device 130(3) which is visible to many users 116. Once selected, the output device 130 may be used to present the user interface 132. Continuing the example, the personally identifiable information may be presented on the display device 130(3) of the user device 134.

The inventory management module 316 may be configured to provide information prospectively which has been previously provided responsive to user speech 128 on other occasions. For example, as the user 116 walks into a particular area of the facility 102 they may be presented with information corresponding to a top three set of inquiries made in that area.

The process for this prospective presentation may include the following. Response data resulting from queries previously generated responsive to speech uttered within a designated area of the facility 102 may be retrieved. For example, response data comprising results of queries responsive to the user speech 128 which has emanated from the users 116 within aisle 112(1) is retrieved from the data store 326.

The inventory management module 316 may determine the first user 116(1) is present in the designated area. Continuing the example, the position data 336 for the first user 116(1) indicates the first user 116(1) is within the aisle 112(1).

The response data may be presented by way of the user interface 132. For example, the response data may be used to generate user interface data 354 which is provided to the output devices 130 in the aisle 112(1). The first user 116(1) has now been presented with the information. The user speech 128 uttered by the first user 116(1) after this presentation may seek additional information, resulting in an additional query. This additional query may be executed to retrieve further information. For example, the user speech 128 may indicate that the first user 116(1) would like to have more detail. This information may be retrieved and presented with the corresponding modification of the user interface 132 based at least in part on the information.

The block 1108 may also be configured to perform actions which include requesting permission from the first user 116(1) to take other actions. In this fashion, the speech interaction module 320 may seek confirmation prior to taking other actions. The first user 116(1) may be presented with the user interface 132 which prompts the first user 116(1) for permission to perform an action. For example, the user interface 132 may comprise synthesized speech asking the first user 116(1) "would you like more information on widgets?" Subsequent performance of the action may be responsive to receiving second user speech 128(2) which is indicative of assent to perform the action. Continuing the example, the first user 116(1) may respond with the user speech 128(2) "yes". Following this, the first user 116(1) may be presented with the user interface 132 which includes the additional information on widgets. As described above, in other implementations the user interfaces 132 may be visual, haptic, audible, and so forth.

The actions performed by the inventory management module 316 may also include dispatching a second user 116(2), a tote 118, robots, or other resource to the first user 116(1). For example, the audio acquired from the first user 116(1) may be processed to determine the presence of non-lexical utterances and item data 332 indicating that a particularly heavy item 104 is being picked, and that the first user 116(1) is having difficulty loading the item 104 onto a

tote 118. The corresponding action taken by the speech interaction module 320 may be to dispatch a second user 116(2) to help with loading.

By using the techniques and systems described in this disclosure, users 116 of the facility 102 may be able to perform their tasks more efficiently and effectively. Instead of having to stop and express a particular inquiry or otherwise directly interact with the inventory management system 122, the users 116 are able to go about their business within the facility 102 receiving assistance in the form of the various actions which may be undertaken automatically, or after assent by the user 116.

The processes discussed herein may be implemented in hardware, software, or a combination thereof. In the context of software, the described operations represent computer-executable instructions stored on one or more computer-readable storage media that, when executed by one or more processors, perform the recited operations. Generally, computer-executable instructions include routines, programs, objects, components, data structures, and the like that perform particular functions or implement particular abstract data types. Those having ordinary skill in the art will readily recognize that certain steps or operations illustrated in the figures above may be eliminated, combined, or performed in an alternate order. Any steps or operations may be performed serially or in parallel. Furthermore, the order in which the operations are described is not intended to be construed as a limitation.

Embodiments may be provided as a software program or computer program product including a non-transitory computer-readable storage medium having stored thereon instructions (in compressed or uncompressed form) that may be used to program a computer (or other electronic device) to perform the processes or methods described herein. The computer-readable storage medium may be one or more of an electronic storage medium, a magnetic storage medium, an optical storage medium, a quantum storage medium, and so forth. For example, the computer-readable storage media may include, but is not limited to, hard drives, floppy diskettes, optical disks, read-only memories (ROMs), random access memories (RAMs), erasable programmable ROMs (EPROMs), electrically erasable programmable ROMs (EEPROMs), flash memory, magnetic or optical cards, solid-state memory devices, or other types of physical media suitable for storing electronic instructions. Further, embodiments may also be provided as a computer program product including a transitory machine-readable signal (in compressed or uncompressed form). Examples of machine-readable signals, whether modulated using a carrier or unmodulated, include but are not limited to signals that a computer system or machine hosting or running a computer program can be configured to access, including signals transferred by one or more networks. For example, the transitory machine-readable signal may comprise transmission of software by the Internet.

Separate instances of these programs can be executed on or distributed across any number of separate computer systems. Thus, although certain steps have been described as being performed by certain devices, software programs, processes, or entities, this need not be the case and a variety of alternative implementations will be understood by those having ordinary skill in the art.

Additionally, those having ordinary skill in the art readily recognize that the techniques described above can be utilized in a variety of devices, environments, and situations. Although the subject matter has been described in language specific to structural features or methodological acts, it is to

39

be understood that the subject matter defined in the appended claims is not necessarily limited to the specific features or acts described. Rather, the specific features and acts are disclosed as exemplary forms of implementing the claims.

The invention claimed is:

1. A system for presenting information in a materials handling facility (facility), the system comprising:

a plurality of microphones configured to continuously acquire audio data representative of sound within the facility;

a plurality of imaging sensors configured to acquire images within the facility;

an output device;

a memory, storing computer-executable instructions; and a hardware processor in communication with the plurality of microphones, the plurality of imaging sensors, the output device, and the memory, wherein the hardware processor is configured to execute the computer-executable instructions to:

identify a first user and a second user using the images acquired within the facility;

determine a relationship between the first user and the second user, wherein the relationship is determined based on the first user and the second user entering the facility at substantially the same time and within a first threshold distance of one another;

process the audio data to detect a portion of the audio data comprising user speech, the user speech comprising natural language spoken by the first user to the second user, wherein the first user and the second user are within a second threshold distance of each other based on the images acquired within the facility;

determine a first user context associated with the first user, wherein the first user context is indicative of one or more of:

the relationship between the first user and the second user,

historical activity of the first user,

a gaze direction of the first user,

position of the first user within the facility,

proximity of the first user to the second user,

demographics of the first user, or

preferences of the first user;

process, using the first user context, the portion of the audio data comprising user speech that is spoken by the first user to the second user to determine one or more key words;

retrieve data responsive to the one or more key words; and

present output representative of the data with the output device.

2. The system of claim 1, wherein the output device comprises a speaker and the output representative of the data comprises speech sounds, and further comprising computer-executable instructions to:

process the audio data to detect second user speech spoken by the first user, wherein the second user speech is indicative of acceptance of the data by the first user; and

store information indicative of the acceptance of the data with the historical activity of the first user; and

wherein the processing of the portion of the audio data to determine the one or more key words comprises computer-executable instructions to:

40

configure operation of a Hidden Markov Model using the first user context.

3. The system of claim 1, further comprising a tote and wherein:

the output device comprises a display device coupled to the tote;

the output comprises a user interface presented on the display device; and further comprising computer-executable instructions to:

determine the first user is speaking to the second user based on one or more of:

facial motions of the first user in the acquired images,

voice recognition, or

sound source localization using the audio data from the plurality of microphones; and

wherein the determination of the first user context is based on the determination the first user is speaking.

4. The system of claim 1, wherein:

the identification of the first user and the second user is based at least in part on one or more of:

facial recognition comprising:

detection of a face in the acquired images,

calculation of one or more eigenvectors of the face, and

comparison of the one or more eigenvectors with previously stored data;

clothing recognition comprising:

detection of one or more articles of clothing in the acquired images,

classification of the one or more articles of clothing, and

comparison of the classifications of the one or more articles of clothing with previously stored data; or

gait recognition comprising:

determination of one or more user limbs in the acquired images,

determination of a change in position over time of the one or more user limbs, and

comparison of the change in position over time with previously stored data; and

wherein the first user context includes identity of the second user;

determine a second user context associated with the second user, wherein the second user context includes the identity of the first user; and

wherein the processing of the portion of the audio data comprising the user speech is further based on the second user context.

5. A computer implemented method comprising:

accessing sensor data obtained from one or more sensors within a material handling facility (facility);

determining a relationship between a first user and a second user, wherein the relationship is based on facility entry data indicative of the first user and the second user entering the facility at substantially the same time and within a first threshold distance of one another;

determining, using the sensor data, a user context associated with a first user of the facility, wherein the user context comprises data indicative of the relationship between the first user and the second user;

acquiring, using one or more microphones in the facility, audio data of user speech generated by the first user while speaking to the second user of the facility in natural language, wherein the first user and the second user are within a second threshold distance of each other based on the sensor data;

41

using the user context, processing the audio data of the user speech with a computing device to determine one or more actions; and
operating, via a network, one or more devices at the facility to perform the one or more actions.

6. The computer implemented method of claim 5, wherein the user context comprises data indicative of one or more items within the facility that the first user is interacting with.

7. The computer implemented method of claim 5, wherein the user context comprises data received from a user device, the data comprising one or more of:

information exchanged using a short message service (SMS),

information exchanged using an instant messaging service,

call metadata associated with an incoming or outgoing telephone or video call, or

application state data indicative of an application executing on the user device.

8. The computer implemented method of claim 5, further comprising:

determining a wordbase of the first user based on one or more of:

word selection in the user speech,

language of the user speech, or

cadence of the user speech; and

wherein the user context further comprises the determined wordbase.

9. The computer implemented method of claim 5, further comprising:

determining an identity of the first user;

accessing user data associated with the identity; and

wherein the user context further comprises the user data.

10. The computer implemented method of claim 9, further comprising:

based at least in part on the user data, generating user interface data;

based at least in part on the user data, selecting an output device; and

wherein the one or more actions comprise:

presenting a user interface by way of the selected output device using the user interface data.

11. The computer implemented method of claim 5, wherein the one or more actions comprise one or more of:

moving an inventory location from a first position in the facility to a second position in the facility, or

moving one or more items from a first inventory location to a second inventory location.

12. The computer implemented method of claim 5, further comprising:

accessing a first proposed route for the first user through the facility;

modifying the first proposed route to generate a second proposed route; and

wherein the one or more actions comprise providing navigational guidance in the facility to the first user based on the second proposed route.

13. The computer implemented method of claim 5, further comprising:

determining a third user is within a third threshold distance of the first user;

generating user interface data; and

sending the user interface data to an output device associated with the first user, wherein the output device is part of a user device or a tote.

42

14. The computer implemented method of claim 5, wherein the natural language does not comprise an inquiry; and

further wherein the one or more actions comprise:

generating a first inquiry from the natural language;

executing a first query to retrieve first information associated with the first inquiry;

executing a second query to retrieve second information from a previous inquiry made by a different user;

generating user interface data using the first information and the second information;

selecting an output device accessible to the first user and other users of the facility; and

processing the user interface data to provide a user interface by way of the selected output device.

15. The computer implemented method of claim 5, wherein the one or more actions comprising:

one or more of:

presenting information designated as personally identifiable information,

presenting information designated as sensitive, or

presenting information about another user;

selecting an output device associated with the first user; and

presenting a user interface by way of the selected output device.

16. The computer implemented method of claim 5, further comprising:

retrieving response data resulting from queries previously generated responsive to speech uttered within a designated area of the facility;

determining the first user is present in the designated area;

presenting by way of a user interface the response data; and

wherein the natural language is processed to generate an inquiry associated with the response data; and

further wherein the one or more actions comprise:

executing a query to retrieve information associated with the inquiry; and

modifying the user interface based on the information.

17. The computer implemented method of claim 5, wherein the one or more actions comprise:

prompting the first user for permission to perform an action; and

wherein the performing the action is responsive to receiving second user speech indicative of assent to perform the action.

18. A computer implemented method comprising:

acquiring, using one or more microphones connected to a network at a material handling facility (facility), audio data representative of sound at the facility;

acquiring, using one or more sensors at the facility that are connected to the network, sensor data;

determining, with one or more computing devices, a relationship between a first user and a second user, wherein the relationship is determined based on the first user and the second user entering the facility at substantially the same time and within a threshold distance of one another;

processing the sensor data using the one or more computing devices to determine a user context of the first user at the facility, wherein the user context is indicative of the relationship between the first user and the second user;

recognizing, with the one or more computing devices, a sound uttered by the first user and directed to the second user at the facility in the acquired audio data;

determining, with the one or more computing devices, an action based on the user context and the sound that is recognized; and

sending from the one or more computing devices, via a network, commands to perform one or more actions 5 associated with operation of the facility.

19. The computer implemented method of claim **18**, wherein the sound is a non-lexical utterance.

20. The computer implemented method of claim **18**, wherein the one or more actions comprise dispatching a 10 third user to assist the first user.

21. The computer implemented method of claim **18**, wherein the one or more actions comprise:

retrieving information;

presenting a first user interface comprising a request for 15 permission to present the retrieved information;

receiving input from the user interface indicative of approval to present the retrieved information; and

presenting a second user interface comprising at least a portion of the information. 20

* * * * *